

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386409
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Dangi; Yogesh et al.

Power-aware scheduling in data centers

Abstract

A computing device determines an expected power consumption value for a set of processes. The computing device further determines an available rack power capacity for one or more server racks amongst multiple server racks, and selects a first server rack having a highest available rack power capacity. The first server rack includes a first set of servers, and the computing device further determines whether a first server of the first set of servers is available. Responsive to determining the first server is available, the computing devices assigns the set of processes to the first server.

Inventors:	Dangi; Yogesh (Pune, IN), Jagadev; Manas Ranjan (San Jose, CA)
Applicant:	NVIDIA Corporation (Santa Clara, CA)
Family ID:	95654595
Assignee:	NVIDIA Corporation (Santa Clara, CA)
Appl. No.:	18/516899
Filed:	November 21, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250165052 A1	May. 22, 2025

Publication Classification

Int. Cl.: G06F1/3228 (20190101)

U.S. Cl.:

CPC G06F1/3228 (20130101);

Field of Classification Search

CPC: G06N (20/00); G06N (7/01); G06F (1/3203); G06F (1/3209); G06F (9/4893); G06F (9/5094); G06F (1/3206); G06F (1/3228); G06F (1/324); G06F (2009/4557); G06F (9/45558); G06F (9/4881); G06F (9/50); G06F (9/5038); G06F (9/505); G06F (1/189); G06F (1/20); G06F (1/206); G06F (1/26); G06F (1/28); G06F (1/3265); G06F (1/3287); G06F (1/329); G06F (13/409); G06F (17/18); G06F (18/214); G06F (2009/45562); G06F (9/5005); G06F (9/5011); G06F (9/5016); G06F (9/5022); G06F (9/5027); G06F (9/5033); G06F (9/5044); G06F (9/5061); G06F (9/5072); G06F (9/5083); G06F (9/5088); G06F (9/522); G06F (9/542); G06F (9/546)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2011/0239010	12/2010	Jain	713/310	G06F 1/3209
2012/0054527	12/2011	Pfeifer	713/340	G06F 1/266
2016/0054781	12/2015	Rajappa	713/320	H04L 47/821
2017/0324813	12/2016	Jain	N/A	G06F 9/5016
2021/0034436	12/2020	Tian	N/A	G06F 9/45558

Primary Examiner: Johnson; Terrell S

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

TECHNICAL FIELD

(1) At least one embodiment pertains to data center process scheduling. For example, at least one embodiment pertains to processors or computing systems used to schedule processes on servers across server racks of a data center.

BACKGROUND

(2) In multi-computing platforms and environments—such as data centers, supercomputers, high-performance computing (HPC) environments, cluster computing environments, or cloud computing environments, etc.—it is important to find idle or underutilized computing devices so that the usages of these computing devices can be more efficiently allocated by taking corrective actions. In the data center or cloud environment, it is important to efficiently use the electrical power provided to a server. When power is not being used, a data center may be underutilizing the computing resources of servers in the data center.

Description

BRIEF DESCRIPTION OF DRAWINGS

(1) Various embodiments in accordance with the present disclosure will be described with reference to the drawings, in which:

(2) FIG. 1A is a block diagram of a data center environment for implementing power-aware

scheduling, according to aspects of the disclosure.

(3) FIG. 1B is a block diagram of a data center environment for implementing power-aware scheduling, according to aspects of the disclosure.

(4) FIG. 1C is a block diagram of a data center environment for implementing power-aware scheduling, according to aspects of the disclosure.

(5) FIG. 1D is a block diagram of a data center environment for implementing power-aware scheduling, according to aspects of the disclosure.

(6) FIG. 2 illustrates a method of power-aware scheduling, according to aspects of the disclosure.

(7) FIG. 3 illustrates a method of power-aware scheduling, according to aspects of the disclosure.

(8) FIG. 4 illustrates a method of power-aware scheduling, according to aspects of the disclosure.

(9) FIG. 5 illustrates a method of power-aware scheduling, according to aspects of the disclosure.

(10) FIG. 6 illustrates an example of a data center, according to aspects of the disclosure.

(11) FIG. 7 is a block diagram illustrating an exemplary computer system which can be a system with interconnected devices and components, a system-on-a-chip (SOC), or some combination thereof, according to aspects of the disclosure.

(12) FIG. 8 is a block diagram illustrating an electronic device for using a processor, according to aspects of the disclosure.

(13) FIG. 9 is a block diagram of a processing system **900**, according to aspects of the disclosure.

DETAILED DESCRIPTION

(14) Embodiments described herein are directed to optimizing data center power usage with power-aware scheduling. A data center can include multiple computing devices. The computing devices can include central processing units (CPUs), graphics processing units (GPUs), data processing units (DPUs), or the like. These computing devices can also be implemented as components in devices referred to as machines, computers, servers, network devices, or the like. These computing devices are important resources in a data center or cloud environment. It is important to have efficient operation of resources in the data center, which can be based on power consumption and/or efficiency of computing devices in the data center. Optimizing available power (electrical energy) consumed by computing devices like servers, is a priority in the data center environment. For example, many data centers have a growing interest in reducing the carbon footprint of the data center, while maintaining mission-critical reliability. In some systems, computing devices can run with a certain fixed power overhead, such that surges in power usage do not cause the computing device (or group of computing devices, such as servers in a server rack) to exceed the electrical power provided to the computing device or group of computing devices. In some systems, the fixed overhead can be large, which can cause computing devices to be underutilized, resulting in a lower efficiency of the data center. In some systems, scheduling jobs can be based on peak power usages by servers which can cause the overhead to be large.

(15) Aspects and embodiments of the present disclosure address these and other challenges by providing a power-aware scheduler for scheduling a set of processes (e.g., applications, jobs, tasks, or routines) received at the data center. By scheduling the set of processes based on an expected power consumption value (e.g., a mean, median, or mode power consumption value for the set of processes), the fixed overhead can be reduced without reducing mission-critical reliability. Additionally, sets of processes can be scheduled based on server rack power to optimize the power usage of a server rack, and not solely scheduled based on optimizing the power usage of individual servers. This can be achieved by causing the power-aware scheduler (e.g., a power-aware scheduling module of a scheduler) to schedule a set of processes on a server of a server rack with the largest available power capacity (e.g., the server with the most electrical power headroom). Additional details regarding determining the largest available power capacity for a server rack (e.g., largest available rack power capacity) are described below with reference to FIGS. 2-3. By scheduling a set of processes on server racks with the most available power, the scheduler can avoid exceeding the peak power usages of the server rack (e.g., the maximum power values

associated with the rack).

(16) Advantages of the disclosure include, but are not limited to, an increased electrical power efficiency for data centers.

(17) FIG. 1A is a block diagram of a data center environment **100A** for implementing power-aware scheduling, according to aspects of the disclosure. The data center environment **100A** includes a data center **101**, set of processes **102**, scheduler **110**, and client devices **104A-N** connected by a network **108**. As described above, the set of processes **102** can include one or more applications, jobs, tasks, routines, or the like.

(18) In at least one embodiment, the data center **101** can include one or more racks **120A-N** and a power monitoring module **130**. Each rack **120** can include one or more multiple computing systems **122A-N** (herein also referred to as “computing system **122**”), where the quantity of racks (M) is a positive integer equal to or greater than zero, and the quantity of computing systems (N) is a second positive integer equal to or greater than zero. In at least one embodiment, each rack **120** can have the same quantity (N) of computing systems **122A-N**. That is, in a data center **101**, the quantity of computing systems **122A-N** can be determined with the equation, $X=M*N$, where X is the quantity of computing system **122A-N** in the data center **101**, M is the quantity of racks **120A-M**, and N is the quantity of computing systems **122A-N** on each rack **120A-M**. In at least one embodiment, data center **101** can refer to a physical location. In at least one embodiment, data center **101** can refer to a logical collection of racks **120A-M**. In at least one embodiment, data center **101** can include additional components, such as network access components, data center services components, etc. In at least one embodiment, network access and/or data center service components can be included in one or more racks **120A-M** and/or one or more computing systems **122A-N** of the data center **101**.

(19) In at least one embodiment, network **108** can include a public network (e.g., the internet), a private network (e.g., a local area network (LAN) or wide area network (WAN)), a wired network (e.g., Ethernet network), a wireless network (e.g., an 802.11 network or a wireless fidelity (Wi-Fi) network), a cellular network (e.g., a Long Term Evolution (LTE) network), routers, hubs, switches, server computers, and/or a combination thereof.

(20) In at least one embodiment, client devices **104A-N** can include a UI dashboard **106**. In at least one embodiment, client devices **104A-N** can generate a set of processes **102** to be received at the scheduler **110** and scheduled on at data center **101** by power-aware scheduling module **112**. In at least one embodiment, client devices **104A-N** can be used to monitor or configure settings of the data center environment **100A** (e.g., through UI dashboard **106**). Client devices **104A-N** can allow users with enhanced privileges with respect to resources of the data center environment **100A** access to components of the data center environment **100A** (e.g., administrator accounts, etc.). In at least one embodiment, the client devices **104A-N** can be used to receive a set of processes **102** from a coupled computing device (e.g., a user device) and forward the set of processes **102** to the scheduler **110**. In at least one embodiment, the client devices **104A-N** can refer to user devices that generates the request including the set of processes that is received at the scheduler **110**.

(21) In at least one embodiment, scheduler **110** includes power-aware scheduling module **112**, processing device **114**, data store **116**, and network connection **118**. The scheduler **110** includes a processing device **114** that can assign a set of processes to a rack **120** and/or computing system **122** of a rack **120** based on a power consumption value for the set of processes, as determined by the power-aware scheduling module **112**. In at least one embodiment, the power consumption value for the set of processes can be predicted by a model, such as a machine learning (ML) model. In at least one embodiment, the ML model can be a random forest regression model. In at least one embodiment, the ML model can be a support vector machine (SVM) model. In at least one embodiment, the ML model can be other types of ML models that can characterize a power consumption value for a given set of processes. In at least one embodiment, the ML model can be trained using historical power consumption values for sets of processes. That is, a set of processes

can be assigned an identification value which can correlate to identification values of previously performed similar sets of processes. In at least one embodiment, the power consumption value can be estimated based on factors associated with the set of processes, such as conditions present in the data center **101**, the rack **120**, and/or the computing system **122**. For example, in at least one embodiment, the set of processes can be received at the power-aware scheduling module **112** along with an estimated power consumption for the set of processes. That is, in at least one embodiment, a system-wide requirement can be imposed to users of client devices **104A-N** (herein also referred to as “client device **104**”) that requires a power consumption value for the set of processes be submitted alongside the set of processes.

(22) In at least one embodiment, the power-aware scheduling module **112** can be coupled to one or more racks **120A-M** and/or one or more computing systems **122A-N** in the data center **101**. Each computing system **122** can include computing resources **124**, such as a CPU, a GPU, a DPV, volatile memory and/or nonvolatile memory. Each computing system **122** can include a network connection **128** to communicate with other devices in the data center **101** and/or other devices over the network **108**. In at least one embodiment, the power-aware scheduling module **112** can be included in the data center **101** (e.g., physically housed in a shared physical location). The power-aware scheduling module **112** can be implemented in one of the computing systems **122**, or as a standalone computing system in the data center **101**. Alternatively, in at least one embodiment, the power-aware scheduling module **112** can be external to the data center **101** (e.g., physically housed in a separate location, such as another data center **101** (not illustrated)). In at least one embodiment, the power-aware scheduling module **112** can schedule sets of processes to be performed across multiple data centers **101** (not illustrated). In at least one embodiment, the scheduler can have one or more critical backup systems (e.g., backup schedulers) which can be configured to perform the functions of power-aware scheduling module **112** if there is an interruption in services from power-aware scheduling module **112**.

(23) In at least one embodiment, the power-aware scheduling module **112** can determine a set of features about a set of processes **102** received at the scheduler **110**. For example, and in at least one embodiment the set of features can include the power consumption value associated with a given set of processes and can retrieve a real-time power usage value of a rack **120** and/or computing system **122** (using power monitoring module **130**). Additional details regarding retrieving real-time power usage values is described below with reference to FIGS. **1B-C**. The power-aware scheduling module **112** can identify the rack **120** with the highest available rack power capacity by subtracting the real-time power usage of the rack from a maximum power value of the rack. For example, if a rack can support 10 kilowatts (e.g., has a maximum rack power value of 10 KW), and the real-time power usage of the rack is 6 KW, then the maximum power value of the rack would be 10 KW-6 KW, or 4 KW. In at least one embodiment, the power-aware scheduling module **112** can collect and store (e.g., in data store **116**) or receive these power usage and power consumption values as a table of values. An example table including illustrative entries, Table 1 is illustrated below:

(24) TABLE-US-00001

TABLE 1	Rack Actual	Rack Available	Percentage	Rack Power	Power										
Usage	Rack Power	of Rack	Identifier	Capacity (Real-Time)	(Capacity – Usage) Power Used										
Rack(a)	8 KW	2 KW	6 KW	25%	Rack(b)	8 KW	4 KW	4 KW	50%		Rack(M)	8 KW	6 KW	2 KW	75%

(25) In at least one embodiment, the power-aware scheduling module **112** can collect power consumption values for given sets of processes over time, and store collected values in data store **116**. In at least one embodiment, power consumption values can be estimated for a given set of processes based on values of entries in data store **116**. Using a processing device (such as processing device **114**) the power-aware scheduling module **112** can estimate a power consumption value for a given set of processes and cause the scheduler **110** to schedule the set of processes on a rack **120** having a highest available rack power capacity. In at least one embodiment, the power-aware scheduling module **112** can send a notification to a user interface (UI) dashboard such as UI

dashboard **106** of a client device **104** to indicate when a set of processes have been assigned to a computing system **122**. In at least one embodiment, the power-aware scheduling module **112** can send a notification to the UI dashboard when the set of processes have been performed. In at least one embodiment, the power-aware scheduling module **112** can send a notification to the UI dashboard **106** indicating the computing system **122A-N** to which a given set of processes have been assigned.

(26) Data store **116** can be a persistent storage that is capable of storing scheduling information. Scheduling information can include information pertaining to the set of processes **102** such as expected power consumption values, estimated power consumption values, actual power consumption values, historical power consumption values, computing resource utilization requirements (e.g., hardware component utilization metrics), duration information, etc. Scheduling information can include information pertaining to the data center **101**, such as available rack power capacities, available server power capacities, real-time rack power usage values, real-time server power usage values, maximum available rack power values, maximum available server power values, etc. Scheduling information can include data structures to tag, organize, and index the scheduling information. Data store **116** can be hosted by one or more storage devices, such as main memory, magnetic or optical storage based disks, tapes or hard drives, network-attached storage (NAS), storage area network (SAN), and so forth. In at least one embodiment, data store **116** can be a network-attached file server, while in other embodiments the data store **116** can be another type of persistent storage such as an object-oriented database, a relational database, and so forth, that can be a separate component of data center environment **100A**, or one or more different machines within the data center environment **100A** (e.g., such as in data center **101**).

(27) In at least one embodiment, the data center environment **100A** can include a system with a memory device and a processing device operatively coupled to the memory device. The computing device can include a set of processing units. The computing device can determine an expected power consumption value for a set of processes. The computing device can determine for racks **120A-N**, an available rack power capacity. The computing device can select from racks **120A-N**, a rack **120** having a highest available rack power capacity. The computing device can determine whether a server of the rack **120** is available. Responsive to determining a server of the rack **120** is available, the computing device can assign the set of processes to the server of the rack **120**. In at least one embodiment, the system includes one or more of a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing deep learning operations; a system for generating synthetic data; a system for generating multi-dimensional assets using a collaborative content platform; a system implemented using an edge device; a system implemented using a robot; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources.

(28) In at least one embodiment, the power-aware scheduling module **112** can collect and store historical power consumption values associated with multiple sets of processes in data store **116** to train a ML model to predict a power consumption value for a newly received set of processes. The ML model can identify various hidden patterns in the historical power consumption values for respective sets of processes and use a current pattern in newly collected data to determine the power consumption value for the newly received set of processes. In at least one embodiment, the ML model can be one or more of a logistics regression model, a k-nearest neighbor model, a random forest regression model, a gradient boost model, or an Extreme Gradient Boost (XGBoost) model. Alternatively, other types of ML models can be used. The trained ML model can be deployed as an object to a component of data center environment **100A** (e.g., such as scheduler **110** or power-aware scheduling module **112**).

(29) FIG. 1B is a block diagram of a data center environment **100B** for implementing power-aware

scheduling, according to aspects of the disclosure. The data center environment **100B** includes a set of processes **102**, the scheduler **110**, the power monitoring module **130**, and racks **120A-N**. For clarity and brevity, not every component of data center environment **100A** is shown and/or described with reference to FIG. **1B**. However, it can be appreciated that FIG. **1B** includes relevant portions of FIG. **1A**, and that insofar as various elements or components are illustrated or described, each component of FIG. **1B** retains the same, or a similar structure to the respectively named and numbered component of FIG. **1A**. That is, for example, scheduler **110** as illustrated in FIG. **1B** can be the same as, or similar to (e.g., perform the functions of) scheduler **110** as illustrated in FIG. **1A**.

(30) A set of processes **102** can be received at scheduler **110**. Power-aware scheduling module **112** can determine a power consumption value for the set of processes, and using the power consumption value for the set of processes **102** and respective maximum power capacities for racks **120A-N**, assign the set of processes **102** to a computing system **122** of a rack **120**. In at least one embodiment, the power-aware scheduling module **112** can obtain maximum power capacities for racks **120A-N** from power monitoring module **130**. In at least one embodiment, the power-aware scheduling module **112** can communicate with the power monitoring module **130** using a HyperText Transfer Protocol (HTTP) Get function. In at least one embodiment, the power-aware scheduling module **112** can assign the set of processes **102** to a rack **120**. In at least one embodiment, power-aware scheduling module **112** can schedule a set of processes (e.g., a “job”) on any computing system **122A-N** (e.g., “server”) of a rack **120** having the highest available rack power capacity.

(31) In at least one embodiment, rack services **129** can be implemented on a computing system **122** of rack **120**, or as a separate computing device. In at least one embodiment, rack services **129** can include a hardware component that can collect real-time power usage values for a rack **120** and/or computing systems **122A-N** of a rack **120**. By way of non-limiting example, in at least one embodiment, the rack services can be implemented as a service, an agent, or a process within the OS or outside the OS in the kernel space of a processing device in a rack **120**, or in scheduler **110** (e.g., such as processing device **114**). In at least one embodiment, rack services **129** can perform one or more functions at regular intervals (e.g., report real-time power usage values, report real-time hardware usage metrics, etc.). In at least one embodiment, rack services **129** can perform one or more functions when triggered, such as when triggered by a power-aware scheduling module **112** attempting to schedule a set of processes **102** at a rack **120** and/or computing system **122** of a rack **120**.

(32) In at least one embodiment, rack services **129** can receive the set of processes **102** and, based on usage metrics associated with computing resources of respective computing systems **122A-N**, assign the set of processes **102** to a computing system **122**. In at least one embodiment, if rack services **129** is unable to assign the set of processes **102** to a computing system **122A-N**, rack services **129** can send the set of processes **102** back to scheduler **110**. For example, the set of processes **102** can more computing resources than are available at any single computing system **122A-N**, or the power consumption value for the set of processes **102** can exceed a maximum available computing system power capacity.

(33) In at least one embodiment, rack services **129** can be used to collect and/or report power usage values for a rack **120** and/or power usage values for respective computing systems **122A-N** of the rack **120**. In at least one embodiment, rack services **129** can be used collect and/or report computing resource usage metrics for a rack **120** and/or computing resource usage metrics for respective computing systems **122A-N** of the rack **120**. Rack services **129** can send power usage values and/or computing resource usage metrics that have been collected for a rack **120** to the power monitoring module **130** and/or the power-aware scheduling module **112** (e.g., to be stored in a data store such as data store **116**). In at least one embodiment, rack services **129** can be performed at each rack **120A-M** (e.g., as a process on one of computing systems **122A-N**). In at least one

embodiment, some of the rack services **129** can be performed as a part of a standalone computing system within the data center environment **100B**. In at least one embodiment, some of the rack services **129** can be performed by the power-aware scheduling module **112** or the scheduler **110**. In at least one embodiment, some of the rack services **129** can be performed by the power monitoring module **130**.

(34) FIG. **1C** is a block diagram of a data center environment **100C** for implementing power-aware scheduling, according to aspects of the disclosure. The data center environment **100C** includes data center **101**, power monitoring module **130**, and representational state transfer (REST) application programming interface (API) services, such as REST API services **138**. A REST API (also known as RESTful API) is an API or web interface that allows interaction with RESTful web services. In at least one embodiment, the REST API services **138** can be exposed to users (e.g., administrators of the data center **101**), where the user can pass feature attributes for a specified period (e.g., an hour/day) as input. The REST API services **138** can provide indications of power usage values of racks **120A-N**, computing systems **122A-N**, and/or power consumption values for a set of processes **102**. As described above, aggregated data can be collected, and used as inputs to the REST API services **138** to predict a power consumption value for a set of processes **102** received at a power-aware scheduling module **112**. In at least one embodiment, information collected, or transmitted using the REST API services **138** can be used to provide a visualization of the data center environment **100C** to a UI dashboard, such as UI dashboard **106** of FIG. **1A**.

(35) For clarity and brevity, not every component of data center environments **100A** and **100B** are shown and/or described with reference to FIG. **1C**. However, it can be appreciated that FIG. **1C** includes relevant portions of FIGS. **1A-B**, and that insofar as various elements or components are illustrated or described, each component of FIG. **1C** retains the same, or similar structure to the respectively named and numbered component of FIGS. **1A-B**. That is, for example, power monitoring module **130** as illustrated in FIG. **1C** can be the same as, or similar to (e.g., perform the functions of) power monitoring module **130** as illustrated and described in FIGS. **1A-B**.

(36) Power monitoring module **130** can include power collection service **132** and data store **134**. Power collection service **132** can collect power usage values from devices and/or components of the data center **101**, such as a rack **120**, or computing system **122** (not illustrated). In at least one embodiment, power collection service **132** can collect power consumption values associated with given sets of processes, and/or historical sets of processes (e.g., sets of processes that have previously been performed). In at least one embodiment, power collection service **132** can retrieve one or more power usage or power consumption values from the data center using the intelligent platform management interface (IPMI) protocol. For example, power collection service **132** can use an IPMI call to fetch real-time (e.g., “current”) power usage values for each computing system **122** (e.g., “server,” not illustrated) in the data center **101**. Once power collection service **132** has obtained power usage values and/or power consumption values, the collected values can be stored in data store **134**. In at least one embodiment, data store **134** can transmit or synchronize the values of stored entries in data store **134** with other data stores in the data center **101** or data center environment **100C** (e.g., such as data store **116** as described with respect to FIG. **1A**). In an illustrative example, and in at least one embodiment, the following IPMI command and power deliver unit (PDU) command can be used to fetch power information from the servers of each rack:

(37) TABLE-US-00002 ipmi-dcml --session-timeout=6000 -D LAN_2_0 -h HOST_IP -u USER_ID -p PASSWORD --get-system-power-statistics snmpget -v 2c -c public PDU_IP FIXED_SUFFIX_OEM PORT

(38) Once real-time power usage values for servers in a data center **101** (e.g., computing systems **122**, not illustrated) have been collected, power monitoring module **130** can obtain server details about each respective server using REST API services **138**. In at least one embodiment, data store **134** of power monitoring module **130** can store a rack-to-server mapping table, that power collection service **132** can use to identify servers (e.g., computing units **122A-N**) of a server rack

(e.g., a rack **120**). An example table including illustrative entries, Table 2 is illustrated below:

(39) TABLE-US-00003

TABLE 2	Rack Identifier	First Server Address	... N-th Server Address
Rack(a)	Server(a)_HOST_IP_(1)	... Server(N)_HOST_IP_(1)	... Rack(M)
	Server(a)_HOST_IP_(M)	... Server(N)_HOST_IP_(M)	

(40) By way of non-limiting example, and in at least one embodiment, power monitoring module **130** can use an IPMI call to retrieve a set of power usage values for a respective set of racks **120A-N** in the data center **101**. The IPMI call can return the power usage values for each rack **120** and a rack identifier. The power monitoring module **130** can store power usage values for each server in a rack **120** of a data center as a table of values in data store **134**. An example table including illustrative entries, Table 3 is illustrated below:

(41) TABLE-US-00004

TABLE 3	Server	Server Power	Actual Server Power	Available Server Power	Percentage of Identifier Capacity Usage (Real Time)	(Capacity – Usage)	Server Power Used		
Server(a)	300 watts	200 watts	100 watts	67%	Server(b)	300 watts	210 watts	90 watts	70%
...	...	...	...	...	...	...	...	...	...
Server(N)	300 watts	240 watts	60 watts	80%					

(42) Using the rack identifier, the power monitoring module **130** can use the REST API services **138** to retrieve rack information from the identified rack. Rack information can include, for example, a maximum rack power value, a rack location (physically and/or logically) server information, server utilization metrics, server power usage values, etc. In at least one embodiment, the power monitoring module **130** can implement a communication using REST API services **138**, which communication is received by another component of data center environment **100C**, such as a power-aware scheduling module **112** (not illustrated). In at least one embodiment, power monitoring module **130** can use the REST API services **138** to fetch the rack information of a rack **120** having the highest available rack power capacity of the racks **120A-N** (not illustrated). In an illustrative example, and in at least one embodiment, the following pseudo code can be used in the REST API services **138**:

(43) TABLE-US-00005

```

{ rack_id: { { "server_ip": "ip_address", "hostname":
"hostname_of_server", "available_cpu": "available_cpu", "available_memory":
"available_memory", }, { "server_ip": "ip_address", "hostname":
"hostname_of_server", "available_cpu": "available_cpu", "available_memory":
"available_memory", }, { "server_ip": "ip_address", "hostname":
"hostname_of_server", "available_cpu": "available_cpu", "available_memory":
"available_memory", } } }

```

(44) The power monitoring module **130** can aggregate power consumption and/or power usage values for a specified time period and summarize the usage data in a summary table, which can be stored in a data store, such as data store **116** of FIG. 1A. This data can be collected and stored for racks, and for servers. Such summary tables can be provided by the REST API service **138**. Exemplary tables including illustrative entries, Table 4 (rack information) and Table 5 (server information) are illustrated below:

(45) TABLE-US-00006

TABLE 4	Rack	Real-Time Rack Power	Available Rack Power	Percentage of Rack Power Used	Rack Identifier
Rack(a)	2 KW	6 KW	25%	YYYY-MM-DD 23:59:59	
...	...	...	...	...	
Rack(M)	6 KW	2 KW	75%	YYYY-MM-DD 23:59:59	

(46) TABLE-US-00007

TABLE 5	Server	Real-Time Server Power	Available Server Power	Percentage of Identifier Power Usage	Server Identifier
Server(a)	200 watts	100 watts	67%	YYYY-MM-DD 23:59:59	
...	...	...	...	...	
Server(N)	240 watts	60 watts	80%	YYYY-MM-DD 23:59:59	

(47) FIG. 1D is a block diagram of a data center environment **100D** for implementing power-aware scheduling, according to aspects of the disclosure. The data center environment **100D** includes set of processes **102**, power-aware scheduling module **112**, and racks **120A-N**. For clarity and brevity, not every component of data center environments **100A-C** are shown and/or described with reference to FIG. 1D. However, it can be appreciated that FIG. 1D includes relevant portions of

FIGS. 1A-C, and that insofar as various elements or components are illustrated or described, each component of FIG. 1D retains the same, or similar structure to the respectively named and number component of FIGS. 1A-C. That is, for example, power-aware scheduling module **112** can be the same as or similar to (e.g., perform the functions of) power-aware scheduling module **112** as illustrated and described in FIGS. 1A-C.

(48) The set of processes **102** can be received at the power-aware scheduling module **112**. The power-aware scheduling module **112** can assign the set of processes to a rack **120A-M** based on various factors, such as a rack maximum power value, the rack real-time power usage, a server maximum power value, a server real-time power usage, and/or available compute resources on a server of a rack **120A-M** (e.g., a computing system **122**, not illustrated). In at least one embodiment, power-aware scheduling module can directly assign the set of processes **102** to a server of a rack **120**. In at least one embodiment, as described with reference to FIG. 1B, the power-aware scheduling module can assign the set of processes **102** to a rack **120**, and rack services (e.g., rack services **129**, not illustrated) can assign the set of processes **102** to a server (e.g., computing system **122**, not illustrated).

(49) In at least one embodiment, the power-aware scheduling module **112** can attempt to identify an available server on a rack **120**, such as the rack **120** having the highest available power capacity (as described with reference to FIGS. 1A-C). The power-aware scheduling module **112** can be unable to find an available server on the rack **120**. The power-aware scheduling module **112** can then identify (e.g., using power monitoring module **130**) another rack **120** with a next-highest available power capacity (e.g., a second-highest available power capacity). In an illustrative example, for a set of three racks **120A-C**, rack **120A** having 2 KW of available power capacity, rack **120B** having 10 KW of available power capacity, and rack **120C** having 5 KW of available power capacity, the power-aware scheduling module **112** can first attempt to identify an available server on rack **120B** (which has the largest available power capacity). If the power-aware scheduling module **112** is unable to identify an available server on rack **120B**, the power-aware scheduling module **112** can next attempt to identify an available server on rack **120C** (the rack **120** with the next-largest available power capacity). If the power-aware scheduling module **112** is unable to identify an available server on rack **120C**, the power-aware scheduling module **112** can next attempt to identify an available server on rack **120A** (the rack **120** with the next-largest available power capacity after rack **120C**, i.e., the rack with the third-largest available power capacity for racks **120A-C**). In at least one embodiment, upon failing to identify an available server on any of racks **120A-N**, power-aware scheduling module **112** can place the set of processes **102** in a waiting queue. In at least one embodiment, upon failing to identify an available server on any of racks **120A-N**, the power-aware scheduling module **112** can request an updated real-time power usage value for each of racks **120A-N** from the power monitoring module **130** (not illustrated).

(50) FIG. 2 illustrates a method **200** of power-aware scheduling, according to aspects of the disclosure. Method **200** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. The processing logic can be implemented in one or more computing devices. In at least one embodiment, method **200** can be performed by power-aware scheduling module **112** of FIGS. 1A-D. In another embodiment, the method **200** can be performed by the processing device **114** of FIG. 1A.

(51) At operation **201**, the processing logic begins the method **200** by identifying a power consumption value for a set of processes. The set of processes can include one or more applications, jobs, tasks, routines, or the like. In at least one embodiment, the power consumption value for the set of processes can be determined using metadata included along with the set of processes. For example, and in at least one embodiment, a power consumption value can be included as metadata along with the set of processes when sent to the scheduler (such as the power-aware scheduling module **112** described with reference to FIGS. 1A-D). In at least one embodiment, the power consumption value for the set of processes can be determined based on the

estimated compute resources of a server (e.g., a computing system **122** described with reference to FIGS. 1A-D) needed to perform the set of processes. Additional details regarding predicting an expected power consumption value for the set of processes based on compute resources of a server is described with reference to FIG. 4.

(52) In at least one embodiment, the power consumption value for the set of processes can be determined based on historical power consumption values for previous sets of processes, or previous similar sets of processes performed at the data center, such as data center **101** described with reference to FIGS. 1A-D. In at least one embodiment, the power consumption value can be predicted a ML model trained on the historical power consumption values. Additional details regarding predicting an expected power consumption value for the set of processes based on historical power consumption values is described with reference to FIG. 5.

(53) At operation **202**, the processing logic selects a first server rack with a highest available rack power capacity. The available rack power capacity can refer to a difference between a real-time power usage value for a given rack and a maximum power value for the rack. For example, referring to Table 1 above, Rack(a) has a rack power capacity of 8 KWs. The real-time power usage of Rack(a) is 2 KWs. Thus, Rack(a) has an available power capacity of 6 KWs. Continuing to refer to Table 1, as illustrated in the “Available Rack Power” column, Rack(a) has the largest (e.g., “highest”) available rack power capacity. Thus, in the illustrative set of Rack(a)-Rack(M), at operation **202**, processing logic would select Rack(a) as the first server rack with the highest available rack power capacity.

(54) At operation **203**, the processing logic determines whether the first server rack has an available server. If the first server rack has an available server, processing logic proceeds to operation **204**. If the first server rack does not have an available server, processing logic proceeds to operation **205**. In at least one embodiment, processing logic can determine whether a server of a server rack is available based on the power usage of the server. Additional details regarding determining whether the server is available based on the power usage of the server are described with reference to FIGS. 3-4. In at least one embodiment, processing logic can determine whether a server of a server rack is available based on usage metrics of hardware components of the server. Additional details regarding determining whether the server is available based on usage metrics of the hardware components of the server are described with reference to FIG. 4.

(55) At operation **204**, responsive to determining the first server rack has an available server, processing logic assigns the set of processes to the available server on the first server rack. In at least one embodiment, processing logic can check whether a first server of the first server rack is available. Responsive to determining the first server of the first server rack is not available, processing logic can check whether a second server of the first server rack is available, etc. Additional details regarding selecting an available server from a rack are described with reference to FIGS. 3-5.

(56) At operation **205**, responsive to determining the first server rack does not have an available server, processing logic identifies a next server rack with a next-highest available rack power capacity. For example, referring to Table 1, as illustrated in the “Available Rack Power” column, Rack(a) has the largest (e.g., “highest”) available rack power capacity, Rack(b) has the second largest (e.g., “next-highest”) available rack power capacity, and Rack(M) has the smallest available rack power capacity. Thus, in the illustrative set of Rack(a)-Rack(M), at operation **205**, processing logic would select Rack(b) as the next server rack with the next-highest available rack power capacity.

(57) In at least one embodiment, where processing logic has started operation **205** in response to failing the operation **206**, processing logic can then select the next server rack with the next-highest available rack power capacity. For example, referring again to Table 1, as illustrated in the “Available Rack Power” column, after processing logic determines that Rack(a) does not have an available server, and has determined that Rack(b) does not have an available server (e.g., the rack

with the “next-highest” available power capacity), then in the illustrative example, the next-highest available power capacity is the available power capacity of Rack(M). Thus, in the illustrative set of Rack(a)-Rack(M), where Rack(a) and Rack(b) have a larger available rack power capacity than Rack(M), but neither Rack(a) nor Rack(b) have an available server, at operation **205**, processing logic would select Rack(M) as the next server rack with the next-highest available rack power capacity. Additional details regarding selecting a server rack with an available rack power capacity (e.g., a “highest,” or “next-highest rack power capacity” are described with reference to FIGS. 3-5).

(58) At operation **206**, processing logic determines whether the next server rack has an available server. In at least one embodiment, processing logic can check whether a first server of the first server rack is available. Responsive to determining the first server of the first server rack is not available, processing logic can check whether a second server of the first server rack is available, etc. Additional details regarding selecting an available server from a rack are described with reference to FIGS. 3-5. If the next server rack does not have an available server, processing logic returns to operation **205**. If the next server rack does have an available server, processing logic proceeds to operation **207**.

(59) At operation **207**, responsive to determining the next server rack has an available server, processing logic assigns the set of processes to the available server on the next server rack.

(60) FIG. 3 illustrates a method **300** of power-aware scheduling, according to aspects of the disclosure. Method **300** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. The processing logic can be implemented in one or more computing devices. In at least one embodiment, method **300** can be performed by power-aware scheduling module **112** of FIGS. 1A-D. In another embodiment, the method **300** can be performed by the processing device **114** of FIG. 1A.

(61) At operation **301**, processing logic begins the method **300** by identifying a power consumption value for a set of processes. The set of processes can include one or more applications, jobs, tasks, routines, or the like. In at least one embodiment, the power consumption value for the set of processes can be determined using metadata included along with the set of processes. For example, and in at least one embodiment, a power consumption value can be included as metadata along with the set of processes when sent to the scheduler (such as the power-aware scheduling module **112** described with reference to FIGS. 1A-D). In at least one embodiment, the power consumption value for the set of processes can be determined based on the estimated compute resources of a server (e.g., a computing system **122** described with reference to FIGS. 1A-D) needed to perform the set of processes. Additional details regarding predicting an expected power consumption value for the set of processes based on compute resources of a server is described with reference to FIG. 4.

(62) In at least one embodiment, the power consumption value for the set of processes can be determined based on historical power consumption values for previous sets of processes, or previous similar sets of processes performed at the data center, such as data center **101** described with reference to FIGS. 1A-D. In at least one embodiment, the power consumption value can be predicted a ML model trained on the historical power consumption values. Additional details regarding predicting an expected power consumption value for the set of processes based on historical power consumption values is described with reference to FIG. 5.

(63) At operation **302**, processing logic obtains a real-time power consumption value for a server rack. In at least one embodiment, processing logic can obtain respective real-time power consumption values for respective server racks of a data center, such as data center **101** as described in reference to FIGS. 1A-D. In at least one embodiment, the real-time power consumption value for a server rack can be obtained with a power monitoring module **130**. The power monitoring module can include hardware, software, and/or firmware and can couple to power deliver units (PDUs) of one or more server racks. In at least one embodiment, a server rack can include multiple PDUs that provide electrical power to multiple servers on the server rack. In at

least one embodiment, the power monitoring module, such as power monitoring module **130** can send a request, such as a IPMI request to one or more PDUs of a server rack that are coupled to servers of the server rack. In response, the power monitoring module can receive a metadata about the one or more PDUs of the server rack, which can include, for example, a real-time power usage value for the server rack. For example, if a server rack includes three PDUs, the IPMI response can indicate the real-time power output of each of the three PDUs. The power monitoring module can then sum the real-time power output of the three PDUs and represent the sum as a rack power usage value.

(64) At operation **303**, processing logic determines whether the server rack has an available power capacity to perform the set of processes based on (i) the power consumption value and (ii) the real-time rack power usage value. If the server rack does not have an available power capacity to perform the set of processes, processing logic returns to operation **302**. If the server rack does have the available power capacity to perform the set of processes, processing logic proceeds to operation **304**. In at least one embodiment, the available power capacity can be calculated by subtracting the real-time rack power usage value from a maximum rack power value. In at least one embodiment, the maximum rack power value can represent a usable power capacity of the rack, with a built-in overhead protection. The ratio of the usable power capacity of the rack to the maximum rack power value can be referred to as the “load factor ratio.” For example, referring to Table 1 above, Rack(a) has a rack power capacity of 8 KW, which represents the usable power of the rack (e.g., the “maximum rack power value”). The power supplied by PDUs of the Rack(a) in this illustrative example can be 10 KW, thus representing a built-in overhead protection of 2 KW, or 20%, and the load-factor ratio is 80%. In this illustrative example, a 20% built-in overhead protection is considered, however, it can be appreciated that the built-in overhead protection can represent another percentage, such as 15%, or 10% (with load factor ratios representing 80%, 85%, and 90% respectively).

(65) At operation **304**, responsive to determining the server rack has an available power capacity to perform the set of processes, processing logic selects a server from the server rack. In at least one embodiment, processing logic can select a server of the server rack based on the power usage of the server. In at least one embodiment, processing logic can select a server based on usage metrics of hardware components of the server. Additional details regarding selecting a server are described with reference to FIG. 4.

(66) At operation **305**, processing logic obtains a real-time power usage value for the server. In at least one embodiment, a server rack can include multiple PDUs that provide electrical power to multiple servers on the server rack. In at least one embodiment, a server can be electrically coupled to two or more PDUs. In at least one embodiment, a server can obtain electrical power from two or more PDUs coupled to the server. For example, a power management module of the data center (which can be included in, or coupled to the power monitoring module **130** described in FIGS. 1A-D) can cause a server to draw 50% of the used electrical power from a first PDU coupled to the server and to draw 50% of the used electrical power from a second PDU coupled to the server. In at least one embodiment, the power monitoring module, such as power monitoring module **130** can send a request, such as a IPMI request to one or more PDUs coupled to servers of a server rack. In response, the power monitoring module can receive a metadata about the PDU, which can include, for example, a real-time power usage value for one or more servers coupled to the PDU.

(67) At operation **306**, processing logic determines whether the server has an available power capacity to perform the set of processes based on (i) the power consumption value and (ii) the real-time server power usage value. If the server does not have an available power capacity to perform the set of processes, processing logic returns to operation **304**. If the server does have an available power capacity to perform the set of processes, processing logic proceeds to operation **307**. In at least one embodiment, the available power capacity can be calculated by subtracting the real-time rack power usage value from a maximum power value associated with a server. In at least one

embodiment, the maximum power value can represent a usable power capacity of the server, with a built-in overhead protection. For example, referring to Table 3 above, Server(a) has a server power capacity of 300 watts, which represents the usable power capacity of the rack (e.g., the “maximum power value.”) The power supplied by PDUs to Server(a) in this illustrative example can be 400 watts, representing a built-in overhead protection of 100 watts, or 25%. In this illustrative example, a 25% built-in overhead protection is considered, however, it can be appreciated that the built-in overhead protection can represent another percentage, such as 10%, or 5% (with usable capacities representing 75%, 90%, and 95% respectively). In at least one embodiment, a server power capacity (e.g., the maximum power value) represents the power supplied by PDUs to the Server. Returning again to Table 3, for example, and in at least one embodiment, the server power capacity of Server(a) (e.g., 300 watts) can represent the power for the server available from one or more PDUs of a server rack, without a protection-overhead tolerance.

(68) At operation **307**, responsive to determining the server has an available power capacity to perform the set of processes, processing logic assigns the set of processes to the server.

(69) FIG. 4 illustrates a method **400** of power-aware scheduling, according to aspects of the disclosure. Method **400** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. The processing logic can be implemented in one or more computing devices. In at least one embodiment, method **400** can be performed by power-aware scheduling module **112** of FIGS. 1A-D. In another embodiment, the method **400** can be performed by the processing device **114** of FIG. 1A.

(70) At operation **401**, the processing logic begins the method **400** by selecting a server rack with a highest available rack power capacity. In at least one embodiment, the real-time power consumption value for a server rack can be obtained with a power monitoring module **130**. The power monitoring module can include hardware, software, and/or firmware and can couple to power deliver units (PDUs) of one or more server racks. In at least one embodiment, a server rack can include multiple PDUs that provide electrical power to multiple servers on the server rack. In at least one embodiment, the power monitoring module, such as power monitoring module **130** can send a request, such as a IPMI request to one or more PDUs of a server rack that are coupled to servers of the server rack. In response, the power monitoring module can receive a metadata about the one or more PDUs of the server rack, which can include, for example, a real-time power usage value for the server rack. For example, if a server rack includes 3 PDUs, the IPMI response can indicate the real-time power output of each of the 3 PDUs. The power monitoring module can then sum the real-time power output of the 3 PDUs and represent the sum as a rack power usage value.

(71) At operation **402**, processing logic selects a server from the server rack. In at least one embodiment, processing logic can select a server of the server rack based on the power usage of the server. In at least one embodiment, processing logic can select a server based on usage metrics of hardware components of the server (e.g., as further described in operation **403**, below).

(72) At operation **403**, processing logic determines whether the server has available compute resources based on usage metrics of hardware components of the selected server. If the server does not have available compute resources, processing logic returns to operation **402**. If the server does have available compute resources, processing logic proceeds to operation **404**. As described above, each server can include various hardware components, such as a CPU, a GPU, a DPU, a volatile memory and/or a nonvolatile memory. A “hardware usage metric” can refer to the ability of the various hardware components to accept a new process, and in at least one embodiment can be represented as a percentage. By way of a non-limiting example, a 75% usage metric for a CPU can represent that the CPU has an available 25% to take on additional processes. Each set of processes can have a required, or estimated compute resource load. That is, a set of processes can be associated with estimated usage metrics of the various hardware components. By way of a non-limiting example, a set of processes can require 5% of a server's CPU, 15% of the server's GPU, 5% of the server's DPU, 30% of the server's volatile memory, and 1% of the server's nonvolatile

memory. Continuing with an illustrative example, if the usage metric corresponding to the server's volatile memory is at 80%, then based on the usage metrics of the hardware components of the server (e.g., the non-volatile memory), the server is unable to perform the set of processes, because $80\%+30\%$ exceeds the available 100%, and thus the server would be unavailable (e.g., the server does not have available compute resources based on usage metrics of the hardware components of the server). Alternatively, continuing with another illustrative example, if the usage metric corresponding to the server's volatile memory is at 40%, and the remaining usage metrics corresponding to hardware components of the server are at 10%, then based on the usage metrics of the hardware components of the server, the server will be able to perform the set of processes, and thus the server would be available (e.g., the server has available compute resources based on usage metrics of the hardware components of the server).

(73) At operation **404**, responsive to determining the server has available compute resources, processing logic identifies a duration for performing a set of processes. In at least one embodiment, the duration for performing the set of processes can be based on the compute resources required to perform the process. In an illustrative example, if a server has a CPU with a certain clock speed, and the set of processes include a certain quantity of operations to be performed by the CPU, the duration for performing the set of processes can be expressed as: $D=Q \cdot R$, where D is a duration expressed as a length of time, Q is the quantity of operations in the set of processes, and R is the quantity of operations that can be performed relative to each clock cycle, based on the CPU clock speed. For example, and in at least one embodiment, the CPU of a server can perform one or more operations per clock cycle, depending on the complexity of the operation, the number of cores in the CPU, and/or the number of threads in the CPU. The duration D can likewise be calculated for each hardware component of the server, such as GPUs, DPUs, and/or volatile or non-volatile memory devices. The duration $D_{\text{sub.M}}$ that has the largest value of the durations D corresponding to each hardware component of the server can be identified as the duration for performing the set of processes. In at least one embodiment, additional considerations can adjust the duration $D_{\text{sub.M}}$, such as bus bandwidth between hardware components, cache speed, and/or other signal latency limitations.

(74) At operation **405**, processing logic predicts, based on the duration for performing the set of processes and usage metrics of the hardware components of the server, an expected power consumption value. The duration for performing the set of processes can indicate how long the set of processes will take to perform, and the usage metrics of the hardware components of the server can indicate what percentage of the hardware component will be used to perform the set of processes. In at least one embodiment, the duration can be used by the power-aware scheduling module, such as power-aware scheduling module **112** of FIGS. **1A-1D** to determine or estimate when a rack and/or server power usage values will change (relative to the estimated power consumption value) based on the completion of the set of processes.

(75) In at least one embodiment, the expected power consumption value can be based on the duration for performing the set of processes, and the full amount of compute resources that the set of processes can use at any given time. In at least one embodiment, processing logic can use a power consumption table stored in a data store to determine the power distribution of a server, with respect to the hardware components. In at least one embodiment, to determine the expected power consumption value for the set of processes, processing logic can multiple the utilization metrics for each hardware component associated with the set of processes by the power distributions associated with each hardware component. In an illustrative example, processing logic can identify a duration of 30-minutes for a set of processes. The full amount of compute resources of the CPU that the set of processes can use is 5% (e.g., performing the set of processes will at most raise the CPU utilization metric by 5%), 10% for the GPU, 10% for the DPU, and 10% for memory. The power consumption table can indicate that the CPU uses 30% of the power provided to the server, the GPU uses 40%, the DPU uses 20% and the memory uses 10%. For a server with a maximum

available power capacity of 300 watts, at full hardware utilization metrics (e.g., 100% utilization for each hardware component), the CPU can use 90 watts, the GPU can use 120 watts, the DPU can use 60 watts, and the memory can use 30 watts. In at least one embodiment, determining the expected power consumption value for the set of processes can be expressed as:

$$(76) P_{Expected} = C(P_c) + G(P_g) + D(P_d) + M(P_m)$$

where P.sub.Expected is the expected power consumption value, C is the CPU utilization rate for the set of processes (e.g., 5%), G is the GPU utilization rate (e.g., 10%), D is the DPU utilization rate (e.g., 10%), M is the memory utilization rate (e.g., 10%); and P.sub.c, P.sub.g, P.sub.d, and P.sub.m are the maximum power usable by the CPU, GPU, DPU, and memory respectively. Thus in the illustrative example, the expected power consumption value P.sub.Expected=0.05(90 watts)+0.1(120 watts)+0.1(60 watts)+0.1(30 watts), or 6.6 watts.

(77) Alternatively, in at least one embodiment, additional data can be used to determine the expected power consumption for the set of processes based on hardware utilization metrics of hardware components of the server. In at least one embodiment, the set of processes can be accompanied with metadata indicating an expected power consumption value to perform the set of processes. In at least one embodiment, the power consumption value for a single hardware component of the hardware components of the server can be used to determine the expected power consumption value for the set of processes. For example, if the power consumption value for a GPU to perform the set of processes is known or determined, processing logic can determine an expected power consumption value for the set of processes based on the power consumption value for the GPU.

(78) At operation **406**, processing logic determines whether the server has an available power capacity to perform the set of processes based on the expected power consumption value. If the server does not have the available power capacity, processing logic returns to operation **402**. If the server does have the available power capacity, processing logic proceeds to operation **407**. In at least one embodiment, the available power capacity can be calculated by subtracting the real-time rack power usage value from a maximum power value. In at least one embodiment, the maximum power value can represent a usable power capacity of the server, with a built-in overhead protection. For example, referring to Table 3 above, Server(a) has a server power capacity of 300 watts, which represents the usable power capacity of the rack (e.g., the “maximum power value”). The power supplied by PDUs to Server(a) in this illustrative example can be 400 watts, representing a built-in overhead protection of 100 watts, or 25%. In this illustrative example, a 25% built-in overhead protection is considered, however, it can be appreciated that the built-in overhead protection can represent another percentage, such as 10%, or 5% (with usable capacities representing 75%, 90%, and 95% respectively). In at least one embodiment, a server power capacity (e.g., the maximum power value) represents the power supplied by PDUs to the Server. Returning again to Table 3, for example, and in at least one embodiment, the server power capacity of Server(a) (e.g., 300 watts) can represent the power for the server available from one or more PDUs of a server rack, without a protection-overhead tolerance.

(79) At operation **407**, responsive to determining that the server has the available power capacity to perform the set of processes, processing logic assigns the set of processes to the server.

(80) FIG. 5 illustrates a method **500** of power-aware scheduling, according to aspects of the disclosure. Method **500** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. The processing logic can be implemented in one or more computing devices. In at least one embodiment, method **500** can be performed by power-aware scheduling module **112** of FIGS. 1A-D. In another embodiment, the method **500** can be performed by the processing device **114** of FIG. 1A.

(81) At operation **501**, the processing logic begins the method **500** by receiving a set or processes. The set of processes can include one or more applications, jobs, tasks, routines, or the like. In at least one embodiment, the set of processes can be received at a scheduler, such as scheduler **110** or

power-aware scheduling module **112** as described with reference to FIGS. **1A-D**. In at least one embodiment, metadata associated with the set of processes can be received along with the set of processes. For example, and in at least one embodiment, a power consumption value can be received as metadata along with the set of processes. In another example, and in at least one embodiment, a job ID can be received as metadata along with the set of processes. The job ID can be unique to the set of processes. In at least one embodiment, the job ID can indicate a job family ID. In at least one embodiment, a job family ID can be received as metadata with the set of processes. For example, and in at least one embodiment, a job ID can uniquely identify a particular set of processes from other sets of processes (concurrent or historical) and a job family ID can identify a type or “family” of sets of processes (e.g., a repeated set of processes run at different times, or in different conditions).

(82) At operation **502**, determines whether the set of processes have previously been performed. If the set of processes have not previously been performed, processing logic jumps to operation **505**. If the set of processes have previously been performed, processing logic proceeds to operation **503**. In at least one embodiment, processing logic can use a job ID to determine whether a given set of processes have previously been performed. In at least one embodiment, processing logic can use a job ID and/or a job family ID to determine whether sets of processes similar to the received set of processes have previously been performed. In at least one embodiment, processing logic can use additional metadata associated with the received set of processes to determine whether the set of processes have previously been performed. For example, and in at least one embodiment, associated metadata can include a power consumption value associated with the set of processes, a compute resource requirement for the set of processes, etc.

(83) At operation **503**, responsive to determining that the set of processes have previously been performed, processing logic obtains one or more historical power consumption values for performing the set of processes. In at least one embodiment, the one or more historical power consumption values for performing the set of processes (or a similar set of processes) can be stored in a data store associated with the scheduler (such as data store **116**). In at least one embodiment, while a set of processes is being performed, a power monitoring module, (e.g., power monitoring module **130** described with reference to FIGS. **1A-D**) can record, in real-time at specified intervals (e.g., every 5 seconds, every 5 minutes, every 5 hours, etc.), the power consumption value associated with the set of processes. In at least one embodiment, the power monitoring module **130** can record a maximum power consumption value associated with performing the set of processes. In an illustrative example, a set of processes can be performed over a 30-minute interval. During the 30-minute interval, a maximum power consumption value associated with performing the set of processes can refer to a peak power usage value during the 30-minute interval that was required to perform the set of processes.

(84) In at least one embodiment, the historical power consumption value for a set of processes can be estimated based on a power usage value of a server performing the set of processes (along with other sets of processes) over the duration of time required to perform the set of processes. In an illustrative example, a first set of processes can be performed by a server over a 30-minute duration. Concurrently, one or more additional sets of processes can be performed by the server during the 30-minute duration. The quantity of the one or more additional sets of processes during the 30-minute duration can fluctuate based on each respective set of processes (e.g., a set of processes can start, stop, and/or start and stop over the 30-minute duration). By using timestamps associated with each set of processes (e.g., a start and end time), and real-time power usage values for the server collected at regular intervals during the 30-minute duration, the power monitoring module **130** can estimate a power consumption value associated with actually performing the set of processes, and store that estimated power consumption value in data store **116** as a historical power consumption value associated with a given set of processes.

(85) At operation **504**, processing logic predicts, based on the one or more historical power

consumption values, an expected power consumption value associated with the set of processes. In at least one embodiment, processing logic can determine an average, mean, median, or mode of historical power consumption values for a set of processes in order to predict the estimated power consumption value. In at least one embodiment, an ML model can be trained on the historical power consumption values for the set of processes. In at least one embodiment, the ML model can be trained on metadata associated with performing previous sets of processes. The ML model can identify various hidden patterns in the historical power consumption values associated with a set of processes or group of sets of processes (e.g., a sets of processes with the same job ID or job family ID). In at least one embodiment, the ML model can be trained using historical power consumption values and ground truth data. As described above, in at least one embodiment, the ML model can be one or more of a logistics regression model, a k-nearest neighbor model, a random forest regression model, a gradient boost model, or an XGBoost model. Alternatively, in at least one embodiment, other types of ML models can be used. The trained ML model can be deployed as an object to a computing device operatively coupled to the power monitoring module **130** and/or a power-aware scheduling module (e.g., power-aware scheduling module **112** as described with reference to FIGS. **1A-D**).

(86) At operation **505**, processing logic selects a server from a server rack. In at least one embodiment, processing logic can select a server based on usage metrics of hardware components of the server.

(87) At operation **506**, processing logic determines whether the server has an available power capacity to perform the set of processes based on the expected power consumption value. If the server does not have the available power capacity to perform the set of processes, processing logic returns to operation **505**. If the server does have the available power capacity to perform the set of processes, processing logic proceeds to operation **507**. In at least one embodiment, the available power capacity can be calculated by subtracting the real-time rack power usage value from a maximum power value of a server. In at least one embodiment, the maximum power value can represent a usable power capacity of the server, with a built-in overhead protection. For example, referring to Table 3 above, Server(a) has a server power capacity of 300 watts, which represents the usable power capacity of the server (e.g., the “maximum power value”). The power supplied by PDUs to Server(a) in this illustrative example can be 400 watts, representing a built-in overhead protection of 100 watts, or 25%. In this illustrative example, a 25% built-in overhead protection is considered, however, it can be appreciated that the built-in overhead protection can represent another percentage, such as 10%, or 5% (with usable capacities representing 75%, 90%, and 95% respectively). In at least one embodiment, a server power capacity (e.g., the maximum power value) represents the power supplied by PDUs to the Server. Returning again to Table 3, for example, and in at least one embodiment, the server power capacity of Server(a) (e.g., 300 watts) can represent the power for the server available from one or more PDUs of a server rack, without a protection-overhead tolerance.

(88) At operation **507**, responsive to determining the server has the available power capacity, processing logic assigns the set of processes to the server.

Data Center

(89) FIG. **6** illustrates an example of a data center **600**, according to aspects of the disclosure. In at least one embodiment, data center **600** includes a data center infrastructure layer **602**, a framework layer **604**, a software layer **606**, and an application layer **608**.

(90) In at least one embodiment, as shown in FIG. **6**, data center infrastructure layer **602** can include a resource orchestrator **610**, grouped computing resources **612**, and node computing resources (node C.R.s) **614A(a)-614N(n)**, where “N” represents any whole, positive integer. In at least one embodiment, node C.R.s **614A(a)-614N(n)** can include, but are not limited to, any number of central processing units (CPUs) or other processors (including accelerators, field-programmable gate arrays (FPGAs), graphics processors, etc.), memory devices (e.g., dynamic

read-only memory), storage devices (e.g., solid-state or disk drives), network input/output (NW I/O) devices, network switches, virtual machines (VMs), power modules, and cooling modules, etc. In at least one embodiment, one or more node C.R.s from among node C.R.s **614A(a)-614N(n)** can be a server having one or more of the above-mentioned computing resources.

(91) In at least one embodiment, grouped computing resources **612** can include separate groupings of node C.R.s housed within one or more racks (not illustrated) or many racks housed in data centers at various geographical locations (also not illustrated). Separate groupings of node C.R.s within grouped computing resources **612** can include grouped compute, network, memory, or storage resources that can be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s, including CPUs or processors, can be grouped within one or more racks to provide compute resources to support one or more workloads. In at least one embodiment, one or more racks can also include any number of power modules, cooling modules, and network switches, in any combination.

(92) In at least one embodiment, resource orchestrator **610** can configure or otherwise control one or more node C.R.s **614A(a)-614N(n)** and/or grouped computing resources **612**. In at least one embodiment, the resource orchestrator **610** can include a software design infrastructure (SDI) management entity for data center **600**. In at least one embodiment, the resource orchestrator **610** can include hardware, software, or some combination thereof.

(93) In at least one embodiment, as shown in FIG. 6, framework layer **604** includes a power-aware scheduler **616**, a configuration manager **618**, a resource manager **620**, and a distributed file system **622**. In at least one embodiment, the power aware scheduler can be a scheduler **110** including a power-aware scheduling module **112** as described with reference to FIG. 1A. In at least one embodiment, framework layer **604** can include a framework to support software **624** of software layer **606** and/or one or more application(s) **626** of application layer **608**. In at least one embodiment, support software **624** or application(s) **626** can respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud, and Microsoft Azure. In at least one embodiment, framework layer **604** can be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter “Spark”) that can use a distributed file system **622** for large-scale data processing (e.g., “big data”). In at least one embodiment, power-aware scheduler **616** can include a Spark driver to facilitate scheduling workloads supported by various layers of data center **600**. In at least one embodiment, configuration manager **618** can be capable of configuring different layers, such as software layer **606** and framework layer **604**, including Spark and distributed file system **622**, for supporting large-scale data processing. In at least one embodiment, resource manager **620** can be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system **622** and power-aware scheduler **616**. In at least one embodiment, clustered or grouped computing resources can include grouped computing resources **612** at data center infrastructure layer **602**. In at least one embodiment, resource manager **620** can coordinate with resource orchestrator **610** to manage these mapped or allocated computing resources.

(94) In at least one embodiment, support software **624** included in software layer **606** can include software used by at least portions of node C.R.s **614A(a)-614N(n)**, grouped computing resources **612**, and/or distributed file system **622** of framework layer **604**. The one or more types of software can include, but are not limited to, internet web page search software, email virus scan software, database software, and streaming video content software.

(95) In at least one embodiment, application(s) **626** included in application layer **608** can include one or more types of applications used by at least portions of node C.R.s **614A(a)-614N(n)**, grouped computing resources **612**, and/or distributed file system **622** of framework layer **604**. One or more types of applications can include, but are not limited to, any number of genomics applications, cognitive computing, and a machine learning application, including training or inferencing software, machine learning framework software (e.g., PyTorch, TensorFlow, Caffe,

etc.) or other machine learning applications used in conjunction with one or more embodiments. (96) In at least one embodiment, any of configuration manager **618**, resource manager **620**, and resource orchestrator **610** can implement any number and type of self-modifying actions based on any amount and type of data acquired in any technically feasible fashion. In at least one embodiment, self-modifying actions can relieve a data center operator of data center **600** from making possibly bad configuration decisions and possibly avoiding underutilized and/or poor-performing portions of a data center.

(97) In at least one embodiment, data center **600** can include tools, services, software, or other resources to train one or more machine learning models or predict or infer information using one or more machine learning models according to one or more embodiments described herein. For example, in at least one embodiment, a machine learning model can be trained by calculating weight parameters according to a neural network architecture using software and computing resources described above with respect to data center **600**. In at least one embodiment, trained machine learning models corresponding to one or more neural networks can be used to infer or predict information using resources described above with respect to data center **600** by using weight parameters calculated through one or more training techniques described herein.

(98) In at least one embodiment, data center **600** can use CPUs, application-specific integrated circuits (ASICs), GPUs, FPGAs, or other hardware to perform training and/or inferencing using the above-described resources. Moreover, one or more software and/or hardware resources described above can be configured as a service to allow users to train or perform inferencing of information, such as image recognition, speech recognition, or other artificial intelligence services.

Computer Systems

(99) FIG. 7 is a block diagram illustrating an exemplary computer system, such as computer system **700**, which can be a system with interconnected devices and components, a system-on-a-chip (SOC), or some combination thereof, according to aspects of the disclosure. In at least one embodiment, computer system **700** can include, without limitation, a component, such as a processor **702**, to employ execution units including logic to perform algorithms for process data, in accordance with the present disclosure, such as in the embodiments described herein. In at least one embodiment, computer system **700** can include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, California, although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) can also be used. In at least one embodiment, computer system **700** can execute a version of WINDOWS' operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux, for example), embedded software, and/or graphical user interfaces, can also be used.

(100) Embodiments can be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, internet Protocol devices, digital cameras, personal digital assistants (PDAs), and handheld PCs. In at least one embodiment, embedded applications can include a microcontroller, a digital signal processor (DSP), a system on a chip, network computers (NetPCs), set-top boxes, network hubs, wide area network (WAN) switches, or any other system that can perform one or more instructions in accordance with at least one embodiment.

(101) In at least one embodiment, computer system **700** can include, without limitation, processor **702** that can include, without limitation, one or more execution units **708** to perform operations according to techniques described herein. In at least one embodiment, computer system **700** is a single-processor desktop or server system, but in another embodiment, the computer system **700** can be a multiprocessor system. In at least one embodiment, processor **702** can include, without limitation, a complex instruction set computer (CISC) microprocessor, a reduced instruction set computing (RISC) microprocessor, a very long instruction word (VLIW) microprocessor, a

processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **702** can be coupled to a processor bus **710** that can transmit data signals between processor **702** and other components in computer system **700**.

(102) In at least one embodiment, processor **702** can include, without limitation, a Level 1 (L1) internal cache memory (cache) **704**. In at least one embodiment, processor **702** can have a single internal cache or multiple levels of internal cache. In at least one embodiment, the cache memory can reside external to processor **702**. Other embodiments can also include a combination of both internal and external caches depending on particular implementation and needs. In at least one embodiment, register file **706** can store different types of data in various registers, including and without limitation, integer registers, floating-point registers, status registers, and instruction pointer registers.

(103) In at least one embodiment, an execution unit **708**, including and without limitation, logic to perform integer and floating-point operations, also reside in processor **702**. In at least one embodiment, processor **702** can also include a microcode (ucode) read-only memory (ROM) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **708** can include logic to handle a power-aware scheduling instruction set **709**. In at least one embodiment, by including power-aware scheduling instruction set **709** in an instruction set of a general-purpose processor, such as processor **702**, along with associated circuitry to execute instructions, operations used by many multimedia applications can be performed using packed data in a general-purpose processor, such as processor **702**. In one or more embodiments, many multimedia applications can be accelerated and executed more efficiently by using the full width of a processor's data bus for performing operations on packed data, which can eliminate the need to transfer smaller units of data across the processor's data bus to perform one or more operations one data element at a time.

(104) In at least one embodiment, execution unit **708** can also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **700** can include, without limitation, a memory **716**. In at least one embodiment, memory **716** can be implemented as a Dynamic Random Access Memory (DRAM) device, a Static Random Access Memory (SRAM) device, a flash memory device, or other memory devices. In at least one embodiment, memory **716** can store instruction(s) **718** and/or data **720** represented by data signals that can be executed by processor **702**.

(105) In at least one embodiment, the system logic chip can be coupled to processor bus **710** and memory **716**. In at least one embodiment, the system logic chip can include, without limitation, a memory controller hub (MCH), such as MCH **714**, and processor **702** can communicate with MCH **714** via processor bus **710**. In at least one embodiment, MCH **714** can provide a high bandwidth memory path **715** to memory **716** for instruction and data storage and for storage of graphics commands, data, and textures. In at least one embodiment, MCH **714** can direct data signals between processor **702**, memory **716**, and other components in computer system **700** and bridge data signals between processor bus **710**, memory **716**, and a system I/O **711**. In at least one embodiment, a system logic chip can provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **714** can be coupled to memory **716** through a high bandwidth memory path **715**, and graphics/video card **712** can be coupled to MCH **714** through an Accelerated Graphics Port (AGP) interconnect **713**.

(106) In at least one embodiment, computer system **700** can use the system I/O **711** that is a proprietary hub interface bus to couple the MCH **714** to I/O controller hub (ICH), such as ICH **730**. In at least one embodiment, ICH **730** can provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, a local I/O bus can include, without limitation, a high-speed I/O bus for connecting peripherals to memory **716**, chipset, and processor **702**. Examples can include, without limitation, data storage **722**, a transceiver **724**, a firmware hub (flash BIOS) **726**, a network controller **728**, a legacy I/O controller **732** containing a user input interface **734**, a serial

expansion port **736**, such as Universal Serial Bus (USB), and an audio controller **738**. In at least one embodiment, data storage **722** can include a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage devices.

(107) In at least one embodiment, FIG. **7** illustrates a computer system **700**, which includes interconnected hardware devices or “chips,” whereas, in other embodiments, FIG. **7** can illustrate an exemplary System on a Chip (SoC). In at least one embodiment, devices can be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe), or some combination thereof. In at least one embodiment, one or more components of computer system **700** are interconnected using compute express link (CXL) interconnects.

(108) FIG. **8** is a block diagram illustrating an electronic device **800** for using a processor **802**, according to aspects of the disclosure. In at least one embodiment, electronic device **800** can be, for example, and without limitation, a notebook, a tower server, a rack server, a blade server, a laptop, a desktop, a tablet, a mobile device, a phone, an embedded computer, or any other suitable electronic device.

(109) In at least one embodiment, electronic device **800** can include, without limitation, processor **802** communicatively coupled to any suitable number or kind of components, peripherals, modules, or devices. In at least one embodiment, processor **802** coupled using a bus or interface, such as an I2C bus, a System Management Bus (SMBus), a Low Pin Count (LPC) bus, a Serial Peripheral Interface (SPI), a High Definition Audio (HDA) bus, a Serial Advance Technology Attachment (SATA) bus, a Universal Serial Bus (USB) (including versions 1, 2, and 3), or a Universal Asynchronous Receiver/Transmitter (UART) bus. In at least one embodiment, FIG. **8** illustrates a system, which includes interconnected hardware devices or “chips,” whereas in other embodiments, FIG. **8** can illustrate an exemplary System on a Chip (SoC). In at least one embodiment, devices illustrated in FIG. **8** can be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe), or some combination thereof. In at least one embodiment, one or more components of FIG. **8** are interconnected using compute express link (CXL) interconnects.

(110) In at least one embodiment, FIG. **8** can include a display **810**, a touch screen **812**, a touch pad **814**, a Near Field Communications unit (NFC) **838**, a sensor hub **826**, a thermal sensor **840**, an Express Chipset (EC), such as EC **816**, a Trusted Platform Module (TPM), such as TPM **820**, BIOS/firmware (FW)/flash memory, such as BIOS, FW Flash **808**, a DSP **854**, a memory drive **806** such as a Solid State Disk (SSD) or a Hard Disk Drive (HDD), a wireless local area network unit (WLAN), such as WLAN unit **842**, a Bluetooth unit **844**, a Wireless Wide Area Network unit (WWAN), such as WWAN unit **850**, a Global Positioning System (GPS) **848**, a camera (USB 3.0 camera) **846**, such as a USB 3.0 camera, and/or a Low Power Double Data Rate (LPDDR) memory unit, such as LPDDR5 **804** implemented in, for example, LPDDR5 standard. These components can each be implemented in any suitable manner.

(111) In at least one embodiment, other components can be communicatively coupled to processor **802** through the components discussed above. In at least one embodiment, processor **802** can include a power-aware scheduling module **830**. In at least one embodiment, an accelerometer **828**, Ambient Light Sensor (ALS), such as ALS **832**, compass **834**, and a gyroscope **836** can be communicatively coupled to sensor hub **826**. In at least one embodiment, thermal sensor **840**, a fan **822**, a keyboard **818**, and a touch pad **814** can be communicatively coupled to EC **816**. In at least one embodiment, speakers **858**, headphones **860**, and microphone **862** can be communicatively coupled to an audio unit **856** which can, in turn, be communicatively coupled to DSP **854**. In at least one embodiment, audio unit **856** can include, for example, and without limitation, an audio coder/decoder (codec) and a class-D amplifier. In at least one embodiment, a subscriber identification module (SIM) card, such as SIM **852** can be communicatively coupled to WWAN unit **850**. In at least one embodiment, components such as WLAN unit **842** and Bluetooth unit **844**, as well as WWAN unit **850** can be implemented in a Next Generation Form Factor (NGFF).

(112) FIG. 9 is a block diagram of a processing system **900**, according to aspects of the disclosure. In at least one embodiment, the processing system **900** includes cache memory **902**, register file **904**, processors **906**, graphics processors **908**, memory controller **910**, interface bus **912**, platform controller hub **914**, and power-aware scheduling module **920**. Processing system **900** can be a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **906** or graphics processors **908**. In at least one embodiment, the processing system **900** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices.

(113) In at least one embodiment, the processing system **900** can include, or be incorporated within a server-based gaming platform, a game console, including a game and media console, a mobile gaming console, a handheld game console, or an online game console. In at least one embodiment, the processing system **900** is a mobile phone, smartphone, tablet computing device, or mobile internet device. In at least one embodiment, the processing system **900** can also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, smart eyewear device, augmented reality device, or virtual reality device. In at least one embodiment, the processing system **900** is a television or set-top box device having one or more processors **906** and a graphical interface generated by one or more graphics processors **908**.

(114) In at least one embodiment, one or more processors **906** each include one or more of the processor cores to process instructions which, when executed, perform operations for system and user software. In at least one embodiment, one or more processors **906** and/or one or more graphics processors can be configured to process a portion of the power-aware scheduling (PAS) instruction set, such as PAS instruction set **922**. In at least one embodiment, PAS instruction set **922** can facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). In at least one embodiment, processor cores can each process a different instruction set from PAS instruction set **922**, which can include instructions to facilitate emulation of other instruction sets (not illustrated). In at least one embodiment, processor cores can also include other processing devices, such as a Digital Signal Processor (DSP).

(115) In at least one embodiment, processors **906** includes cache memory **902**. In at least one embodiment, processors **906** can have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory **902** is shared among various components of processors **906**. In at least one embodiment, processors **906** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not illustrated), which can be shared among processor cores using known cache coherency techniques. In at least one embodiment, register file **904** is additionally included in processors **906**, which can include different types of registers for storing different types of data (e.g., integer registers, floating-point registers, status registers, and an instruction pointer register). In at least one embodiment, register file **904** can include general-purpose registers or other registers.

(116) In at least one embodiment, one or more processors **906** are coupled with one or more interface bus **912** to transmit communication signals such as address, data, or control signals between processor cores and other components in processing system **900**. In at least one embodiment, interface bus **912**, in one embodiment, can be a processor bus, such as a version of a Direct Media Interface (DMI) bus. In at least one embodiment, interface bus **912** is not limited to a DMI bus, and can include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI Express), memory busses, or other types of interface busses. In at least one embodiment, processors **906** include an integrated memory controller (e.g., memory controller **910**) and a platform controller hub **914** (PCH). In at least one embodiment, memory controller **910** facilitates communication between a memory device and other components of the processing system **900**, while platform controller hub **914** provides connections to I/O devices via a local I/O bus.

(117) In at least one embodiment, the memory device **930** can be a dynamic random access

memory (DRAM) device, a static random access memory (SRAM) device, a flash memory device, a phase-change memory device, or some other memory device having suitable performance to serve as process memory. In at least one embodiment, the memory device **930** can operate as system memory for processing system **900** to store instructions **932** and data **934** for use when one or more processors **906** executes an application or process. In at least one embodiment, memory controller **910** also optionally couples with an external processor **938**, which can communicate with one or more graphics processors **908** in processors **906** to perform graphics and media operations. In at least one embodiment, a display device **936** can connect to processors **906**. In at least one embodiment, the display device **936** can include one or more of an internal display device, as in a mobile electronic device or a laptop device, or an external display device attached via a display interface (e.g., DisplayPort, etc.). In at least one embodiment, display device **936** can include a head-mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

(118) In at least one embodiment, the platform controller hub **914** enables peripherals to connect to memory device **930** and processors **906** via a high-speed I/O bus. In at least one embodiment, I/O peripherals include, but are not limited to, a data storage device **940** (e.g., hard disk drive, flash memory, etc.), a touch sensor **942**, a wireless transceiver **944**, firmware interface **946**, a network controller **948**, or an audio controller **950**.

(119) In at least one embodiment, the data storage device **940** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI Express). In at least one embodiment, touch sensor **942** can include touch screen sensors, pressure sensors, or fingerprint sensors. In at least one embodiment, wireless transceiver **944** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (LTE) transceiver. In at least one embodiment, firmware interface **946** enables communication with system firmware and can be, for example, a unified extensible firmware interface (UEFI). In at least one embodiment, the network controller **948** can enable a network connection to a wired network. In at least one embodiment, a high-performance network controller (not illustrated) couples with interface bus **912**. In at least one embodiment, audio controller **950** can be a multi-channel high-definition audio controller. In at least one embodiment, the processing system **900** includes an optional legacy I/O controller **952** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to the processing system **900**. In at least one embodiment, the platform controller hub **914** can also connect to one or more Universal Serial Bus (USB) controllers, such as USB controller **960** to connect input devices, such as a keyboard and mouse combination (keyboard/mouse **962**), a camera **964**, or other USB input devices.

(120) In at least one embodiment, an instance of memory controller **910** and platform controller hub **914** can be integrated into a discreet external graphics processor, such as external processor **938**. In at least one embodiment, the platform controller hub **914** and/or memory controller **910** can be external to one or more processors **906**. For example, in at least one embodiment, the processing system **900** can include an external memory controller (e.g., memory controller **910**) and the platform controller hub **914**, which can be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with the processors **906**.

(121) Other variations are within the spirit of the present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit the disclosure to a specific form or forms disclosed, on the contrary, the intention is to cover all modifications, alternative constructions, and equivalents falling within the spirit and scope of the disclosure, as defined in appended claims.

(122) Use of terms “a” and “an” and “the” and similar referents in the context of describing disclosed embodiments (especially in the context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and

not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. The term “connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitations of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within the range, unless otherwise indicated herein, and each separate value is incorporated into the specification as if it were individually recited herein. Use of the term “set” (e.g., “a set of items”) or “subset,” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, the term “subset” of a corresponding set does not necessarily denote a proper subset of the corresponding set, but the subset and corresponding set may be equal.

(123) Conjunctive language, such as phrases of the form “at least one of A, B, and C,” or “at least one of A, B, and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with the context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of a set of A and B and C. For instance, in an illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B, and C” refer to any of the following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain at least one embodiment require at least one of A, at least one of B, and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, the term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). A plurality is at least two items but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, the phrase “based on” means “based at least in part on” and not “based solely on.”

(124) Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under the control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause a computer system to perform operations described herein. A set of non-transitory computer-readable storage media, in at least one embodiment, comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lacks all of the code while multiple non-transitory computer-readable storage media collectively store all of the code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors—for example, a non-transitory computer-readable storage medium stores instructions, and a main central processing unit (CPU) executes some of the instructions while a graphics processing unit (GPU) executes other instructions. In at least one embodiment, different

components of a computer system have separate processors, and different processors execute different subsets of instructions.

(125) Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein, and such computer systems are configured with applicable hardware and/or software that enable the performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

(126) Use of any and all examples or exemplary language (e.g., “such as”) provided herein is intended merely to better illuminate embodiments of the disclosure and does not pose a limitation on the scope of the disclosure unless otherwise claimed. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the disclosure.

(127) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to the same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(128) In description and claims, the terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may not be intended as synonyms for each other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other but yet still co-operate or interact with each other.

(129) Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system or similar electronic computing device, that manipulates and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

(130) In a similar manner, the term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As non-limiting examples, a “processor” may be a CPU or a GPU. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes for carrying out instructions in sequence or in parallel, continuously, or intermittently. The terms “system” and “method” are used herein interchangeably insofar as a system may embody one or more methods, and methods may be considered a system.

(131) In the present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. Obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways, such as by receiving data as a parameter of a function call or a call to an application programming interface. In some implementations, the process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In another implementation, the process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. References may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, the process of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of

an application programming interface, or an interprocess communication mechanism.

(132) Although the discussion above sets forth example implementations of described techniques, other architectures may be used to implement described functionality and are intended to be within the scope of this disclosure. Furthermore, although specific distributions of responsibilities are defined above for purposes of discussion, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

(133) Furthermore, although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. A method comprising: determining, using a computing device, an expected power consumption value for a set of processes that are scheduled to be performed by a server of a server rack of a plurality of server racks; determining an available rack power capacity for one or more server racks of the plurality of server racks; selecting, from the one or more server racks, a first server rack having a highest available rack power capacity, the first server rack comprising a first set of servers; determining whether a first server of the first set of servers is available; and responsive to determining the first server is available, assigning the set of processes to the first server.
2. The method of claim 1, wherein determining, for the one or more server racks of the plurality of server racks, the available rack power capacity comprises: obtaining a real-time rack power usage value for at least one server rack of the one or more server racks; and subtracting the real-time rack power usage value from a maximum available rack power value, wherein the maximum available rack power value is based on a load factor ratio of a maximum rack power value.
3. The method of claim 1, wherein determining whether the first server is available comprises: determining whether the first server has available compute resources; and determining whether the first server has an available power capacity.
4. The method of claim 3, wherein the available compute resources of the first server are based on usage metrics of hardware components of the first server, wherein determining the expected power consumption value comprises: determining a duration for performing the set of processes; and predicting, based on the duration and the usage metrics of the hardware components of the first server, the expected power consumption value.
5. The method of claim 3, wherein determining whether the first server has the available power capacity comprises: obtaining a real-time power usage value for the first server; determining whether a first sum of the real-time power usage value and the expected power consumption value exceeds a maximum power value associated with the first server; and responsive to determining the first sum does not exceed the maximum power value, indicating that the first server has the available power capacity.
6. The method of claim 1, further comprising: responsive to determining the first server is not available, determining whether a second server of the set of servers is available; and responsive to determining the second server is available, assigning the set of processes to the second server.
7. The method of claim 1, further comprising: responsive to determining the first server is not available, selecting, from the one or more server racks, a second server rack having a second-highest available rack power capacity, the second server rack comprising a second set of servers; and determining whether a third server of the second set of servers is available; and responsive to determining the third server is available, assigning the set of processes to the third server.
8. The method of claim 1, wherein determining the expected power consumption value comprises: determining if the set of processes have previously been performed; responsive to determining the set of processes have previously been performed, obtaining from a data store, one or more

historical power consumption values for the set of processes; and predicting the expected power consumption value based on the one or more historical power consumption values.

9. The method of claim 8, wherein predicting the expected power consumption value comprises: determining using a machine learning (ML) model, the expected power consumption value, wherein the ML model is updated using at least of the one or more historical power consumption values.

10. The method of claim 9, wherein the ML model is at least one of a random forest regression model or a support vector machine (SVM) model.

11. A computing device comprising: one or more processing units to determine an expected power consumption value for a set of processes that are scheduled to be performed, determine an available rack power capacity for one or more server racks of a plurality of server racks, select from the one or more server racks, a first server rack having a highest available rack power capacity, the first server rack comprising a first set of servers, and responsive to determining a first server of the first set of servers is available, assign the set of processes to the first server.

12. The computing device of claim 11, wherein to determine an available rack power capacity for at least one server rack of the one or more server racks, the one or more processing units are to: obtain a real-time rack power usage values for the at least one server rack; and subtract the real-time rack power usage value from a maximum available rack power value, wherein the maximum available rack power value is based on a load factor ratio of a maximum rack power value.

13. The computing device of claim 11, wherein to determine whether the first server is available, the one or more processing units are to: determine whether the first server has available compute resources; and determine whether the first server has an available power capacity.

14. The computing device of claim 13, wherein the available compute resources of the first server are based on usage metrics of hardware components of the first server, wherein to determine the expected power consumption value, the one or more processing units are to: determine a duration to perform the set of processes; and predict, based on the duration and the usage metrics of the hardware components of the first server, the expected power consumption value.

15. The computing device of claim 13, wherein to determine whether the first server has the available power capacity, the one or more processing units are to: obtain a real-time server power usage value for the first server; determine whether a first sum of the real-time server power usage value and the expected power consumption value exceeds a maximum power value associated with the first server; and responsive to determining the first sum does not exceed the maximum power value, indicate that the first server has available power capacity.

16. A system comprising: a memory device; and a processing device coupled to the memory device, wherein the processing device is to: determine an expected power consumption value for a set of processes that are schedule to be performed by a server of a server rack of a plurality of server racks; determine an available power capacity for one or more server racks of the plurality of server racks; select, from the one or more server racks, a first server rack having a highest available rack power capacity, wherein the first server rack comprises a first set of servers; determine whether a first server of the first set of servers is available; and responsive to determining the first server is available, assign the set of processes to the first server.

17. The system of claim 16 wherein to determine the expected power consumption value, the processing device is to: determine if the set of processes have previously been performed; responsive to determining the set of processes have previously been performed, obtain from a data store, one or more historical power consumption values for the set of processes; and predict the expected power consumption value based on the one or more historical power consumption values.

18. The system of claim 17, wherein to predict the expected power consumption value, the processing device is to: determine using a machine learning (ML) model, the expected power consumption value, wherein the ML model is updated using at least one of the one or more historical power consumption values.

19. The system of claim 18, wherein the ML model is at least one of a random forest regression model or a support vector machine (SVM) model.

20. The system of claim 16, wherein the system comprises one or more of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing deep learning operations; a system for generating synthetic data; a system for generating multi-dimensional assets using a collaborative content platform; a system implemented using an edge device; a system implemented using a robot; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources.
