

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386537
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Otto; Erich Stephen et al.

Zone segment drive management

Abstract

Techniques for management of data storage in distributed storage systems are provided. A method may include receiving, by a computer system, a request to write data to a volume. The method may include identifying, by the computer system, a zone segment mapped to the volume. The zone segment may include a plurality of zones. The method may include identifying, by the computer system, a segment pointer indicating a write location in a zone of the zone segment. The method may include writing, by the computer system, the data to one or more zones of the plurality of zones of the zone segment, starting at the write location. The method may also include updating, by the computer system, the segment pointer according to a data endpoint of the data in the zone segment.

Inventors: Otto; Erich Stephen (Longmont, CO), Kirby; Michael Warren (Westminister, CO), Kulkarni; Ajinkya Pramod (Thornton, CO)

Applicant: Oracle International Corporation (Redwood Shores, CA)

Family ID: 77499904

Assignee: Oracle International Corporation (Redwood Shores, CA)

Appl. No.: 18/583569

Filed: February 21, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240192876 A1	Jun. 13, 2024

Related U.S. Application Data

continuation parent-doc US 18074206 20221202 US 11947820 child-doc US 18583569
continuation parent-doc US 17229709 20210413 US 11543985 20230103 child-doc US 18074206

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

CPC G06F3/064 (20130101); G06F3/0635 (20130101); G06F3/0653 (20130101); G06F3/0665 (20130101); G06F3/0679 (20130101); G06F3/0604 (20130101)

Field of Classification Search

CPC: G06F (3/0604); G06F (3/0608); G06F (3/061); G06F (3/0635); G06F (3/064); G06F (3/0644); G06F (3/0653); G06F (3/0665); G06F (3/067); G06F (3/0676); G06F (3/0679); G06F (16/1727)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8699185	12/2013	Teh et al.	N/A	N/A
9418699	12/2015	Gibbons et al.	N/A	N/A
9727588	12/2016	Ostapovicz et al.	N/A	N/A
10175892	12/2018	Kim et al.	N/A	N/A
10417190	12/2018	Donlan et al.	N/A	N/A
2011/0197035	12/2010	Na et al.	N/A	N/A
2016/0224260	12/2015	Bandic et al.	N/A	N/A
2017/0168930	12/2016	Choi	N/A	G06F 3/0659
2020/0090684	12/2019	Ikeshima	N/A	N/A
2021/0397618	12/2020	Kim et al.	N/A	N/A
2022/0413706	12/2021	Ye et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
105843753	12/2015	CN	N/A
110910907	12/2019	CN	N/A
111399762	12/2019	CN	N/A
2011-165179	12/2010	JP	N/A
10-2020-0119542	12/2019	KR	N/A

OTHER PUBLICATIONS

International PCT Application No. PCT/US2021/042619, "International Search Report and the Written Opinion", dated Dec. 23, 2021, 16 pages. cited by applicant

U.S. Appl. No. 17/229,709, Non-Final Office Action mailed on Apr. 28, 2022, 32 pages. cited by applicant

U.S. Appl. No. 17/229,709, Notice of Allowance mailed on Sep. 19, 2022, 7 pages. cited by applicant

Primary Examiner: Mackall; Larry T

Background/Summary

INCORPORATION BY REFERENCE; DISCLAIMER

(1) Each of the following applications are hereby incorporated by reference: application Ser. No. 18/074,206, filed on Dec. 2, 2022; application Ser. No. 17/229,709, filed on Apr. 13, 2021. The applicant hereby rescinds any disclaimer of claims scope in the parent application(s) or the prosecution history thereof and advises the USPTO that the claims in the application may be broader than any claim in the parent application(s).

BACKGROUND

(2) Cloud-based platforms provide scalable and flexible computing resources for users' data. Such cloud-based platforms, also referred to as infrastructure as a service (IaaS), may offer entire suites of cloud solutions around a customer's data, for example, solutions for authoring transformations, loading data, and presenting the data.

(3) Shingled Magnetic Recording (SMR) increases storage capacity on hard disk drives by partially overlapping, or shingling, recording tracks. The overlapped tracks may be grouped into zones. By shingling tracks in such a way that a read path is preserved, the tracks may still be read even though they have been partially overwritten. In contrast, as a result of the shingling, while read operations may be still be performed randomly within the zones, write operations will need to be applied sequentially to an entire zone, rather than to a range within the zone.

BRIEF SUMMARY

(4) Techniques are provided (e.g., a method, a system, non-transitory computer-readable medium storing code or instructions executable by one or more processors) for management of object storage of cloud resources data on shingled magnetic recording drives.

(5) In certain embodiments, a method may include receiving, by a computer system, a request to write data to a volume. The method may include identifying, by the computer system, a zone segment mapped to the volume. The zone segment may include a plurality of zones. The method may include identifying, by the computer system, a segment pointer indicating a write location in a zone of the zone segment. The method may include writing, by the computer system, the data to one or more zones of the plurality of zones of the zone segment, starting at the write location. The method may also include updating, by the computer system, the segment pointer according to an endpoint of the data in the zone segment.

(6) In some embodiments, the method may further include ascertaining a size of the data identified in the request. The method may include ascertaining a capacity of the zone, wherein the capacity corresponds to a write path between the write location and an endpoint of the zone. The method may include comparing the capacity of the zone to the size of the data. The method may include, in accordance with the size of the data exceeding the capacity of the zone, breaking the data into a first block and a second block. The size of the first block may correspond to the capacity of the zone. The method may include writing the first block starting at the write location. The method may also include writing the second block to a subsequent zone of the zone segment.

(7) In some embodiments, updating the segment pointer may include, in accordance with the endpoint of the data in the zone segment coinciding with a capacity of the zone segment, marking the zone segment as a full zone segment. Updating the segment pointer may also include identifying a new zone segment mapped to the volume, and updating the segment pointer to the starting point of the new zone segment. Identifying the segment pointer may include scanning the zone segment mapped to the volume for empty or open zones.

(8) In some embodiments, a computer system includes one or more processors and a memory in

communication with the one or more processors, the memory configured to store computer-executable instructions, wherein executing the computer-executable instructions causes the one or more processors to perform one or more of the steps of the method or its variations described above. Identifying the segment pointer may further include, in accordance with the zone segment mapped to the volume not including an empty or open zone, allocating a new zone segment to the volume. The segment pointer may be the start position of a first zone of the new zone segment.

(9) In some embodiments, the method may include ascertaining a usage fraction of the zone segment. The method may include comparing the usage fraction to a usage threshold. The method may include, in accordance with the usage fraction not satisfying the usage threshold, identifying stored data in the plurality of zones of the zone segment, mapping the stored data from one or more source zones to one or more empty or open zones of the plurality of zones, writing the stored data to the empty or open zones, and resetting the source zones to an open status. Identifying stored data in the plurality of zones may include receiving mapping metadata describing the mapping of the zone segment to the volume and locating the stored data in the plurality of zones according to the mapping metadata.

(10) In some embodiments, the plurality of zones may be implemented in a shingled magnetic recording storage system.

(11) In certain embodiments, a computer system includes one or more processors and a non-transitory computer-readable memory storing instructions that, when executed by the one or more processors, cause the one or more processors to perform one or more operations of the method or its variations described above.

(12) In certain embodiments, a computer-readable storage medium stores computer-executable instructions that, when executed, cause one or more processors of a computer system to perform one or more operations of the method or its variations described above.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 illustrates an example system for managing object storage servers, in accordance with one or more embodiments.
- (2) FIG. 2 illustrates an example shingled magnetic recording drive, in accordance with one or more embodiments.
- (3) FIG. 3 illustrates an example flow for writing object data to a zone segment, in accordance with one or more embodiments.
- (4) FIG. 4 illustrates an example flow for reading data from a zone segment, in accordance with one or more embodiments.
- (5) FIG. 5 illustrates an example flow for deleting data from a zone segment, in accordance with one or more embodiments.
- (6) FIG. 6 illustrates an example flow for defragmenting a zone segment, in accordance with one or more embodiments.
- (7) FIG. 7 is a block diagram illustrating one pattern for implementing a cloud infrastructure as a service system, according to at least one embodiment.
- (8) FIG. 8 is a block diagram illustrating another pattern for implementing a cloud infrastructure as a service system, according to at least one embodiment.
- (9) FIG. 9 is a block diagram illustrating another pattern for implementing a cloud infrastructure as a service system, according to at least one embodiment.
- (10) FIG. 10 is a block diagram illustrating another pattern for implementing a cloud infrastructure as a service system, according to at least one embodiment.
- (11) FIG. 11 is a block diagram illustrating an example computer system, according to at least one

embodiment.

DETAILED DESCRIPTION

(12) In the following description, various embodiments will be described. For purposes of explanation, specific configurations and details are set forth in order to provide a thorough understanding of the embodiments. However, it will also be apparent to one skilled in the art that the embodiments may be practiced without the specific details. Furthermore, well-known features may be omitted or simplified in order not to obscure the embodiment being described.

(13) Cloud-based platforms provide scalable and flexible computing resources for users. Such cloud-based platforms, also referred to as infrastructure as a service (IaaS) may offer entire suites of cloud solutions around a customer's data, for example solutions for authoring transformations, loading data, and presenting the data. In data replication and backup systems, which may be a part of a back end data storage system, Shingled Magnetic Recording (SMR) drives may increase storage capacity of the data storage system. SMR drives partially overlapping, or shingling, recording tracks, and thereby may permit more tracks to be written to a hard disk of a given size. The overlapped tracks may be grouped into zones, that may be organized into zone segments. By shingling tracks in such a way that a read path is preserved, the tracks may still be read even though they have been partially overwritten. In contrast, as a result of the shingling, a write operation may be applied in an append-only process, by which partially overwritten data may be read, but may not be overwritten. As such, access to SMR drives may be provided by a block interface, rather than a file system, as described below, due to the additional operations involved in writing to SMR drives.

(14) In distributed storage systems, user data may be stored in volume storage (e.g., block volume storage) and in object storage. Users may interact directly with volume storage systems, while a back-end storage subsystem of the database system may use object storage, for example, for data replication, backup, or other data storage (e.g., to supplement database system capacity).

Implementation of SMR drives in data replication or backup systems is limited by the relative complexity of writing and rewriting data to SMR drives. As described in more detail in reference to FIG. 2, organizing partially overwritten tracks in zones may cause inefficiency when a system attempts to write to a zone data that exceeds the available capacity in the zone. For example, in conventional SMR drives, a system may reject a request to write to a zone data exceeding the capacity of the zone. To address this limitation, zones may be organized into zone segments, and a system managing data storage may address write requests to the zone segment rather than to an individual zone. In this way, a write position in a zone segment may be identified, and data referenced by a write request may be written across multiple zones in a zone segment.

(15) In some examples, block storage is often used where fast, efficient, and reliable data transportation is desired. Block storage breaks up data into blocks, and can store those blocks as separate pieces (e.g., each block with its own identifier). The blocks can then be stored across different systems and each block can be configured (e.g., partitioned) to work with different systems. Alternatively, object storage breaks data file into pieces (e.g., objects), and then stores them in a single repository, which can be spread across multiple networked systems. In the context of block storage, a volume (e.g., a block volume) can be a detachable (e.g., physically or virtually) block storage device that allows a developer to dynamically expand the storage capacity of an instance. To increase or decrease the size of block volume storage, block volumes can be spun up or down, respectively, without much effort. Where the database system implements both volume storage and object storage, I/O requests may make reference to data by a volume identifier (e.g., a block volume identifier). The volume identifier may be mapped to storage locations in the object storage systems, through volume metadata stored in a metadata database included as part of the database system. In the case of a write request, the database system may identify a zone segment in an SMR system storing data mapped to a volume referenced in the write request. Once identified, the database system may identify a zone segment pointer that describes a write position in a zone included in the zone segment at which new data may be written. Rather than overwriting the zone,

the SMR system may write to the zone, may break the data into blocks to be written across multiple zones, may decline the request if a zone is not available having the capacity to store the data, or may allocate new zone segment(s) to store the overrun data. In this way, the SMR system may write data to hard disk drives in such a way that the effective capacity of the object storage system is increased through the use of SMR storage, while limiting the impact of zone write failures caused by data size exceeding zone capacity. In a similar way, reading, deleting, and defragmenting techniques may also be implemented using the SMR system, in reference to volume metadata mapped to data stored in zone segments.

(16) In an illustrative example, a database system may be configured to store cloud resource data as part of a distributed storage system. The database system may include an SMR system. The SMR system may include an SMR storage server in communication with one or more SMR database systems, which may be or include hard disk drives configured for SMR storage using zones organized in zone segments. In this example, the system may receive input/output (I/O) requests, such as write requests (also referred to as “put” requests), read requests (also referred to as “get” requests), as well as other I/O processes including, but not limited to, delete requests or defragmentation operations. As part of SMR configuration, I/O requests may cause the database system to implement improved processes to identify and manage data stored in the SMR system in zone segments.

(17) The process of implementing I/O processes with the back-end storage subsystem improves the overall performance of data replication, storage, and backup systems, at least because it permits efficient I/O from SMR configured storage systems. For example, SMR systems may provide an increased storage capacity of as much as 25%, or more, without addition of hard disk drives to the system, by at least partially shingling previously written magnetic tracks. Organizing SMR zones into zone segments may improve I/O processes by permitting improved speed, efficiency, and success rate of serving I/O requests made on SMR drives. Furthermore, by managing zone drives as described in reference to the forthcoming figures, the database system may reduce system complexity by permitting the database system to use host-managed drives, controlled by an integrated block interface making reference to volume metadata to map object data to volume data.

(18) FIG. 1 illustrates an example system **100** for managing object storage servers, in accordance with one or more embodiments. The database system **100** may be configured to distribute and store user data across multiple storage systems, which may implement various magnetic recording technologies (e.g., perpendicular magnetic recording (PMR), SMR, or the like), as part of a back end data storage subsystem. In this way, the database system **100** may serve I/O requests on data stored in object storage as part of data replication, restoring volume data to block volume systems, or defragmenting SMR systems that organize zones into zone segments as an approach to facilitating improved I/O operations in a distributed data storage system.

(19) In some embodiments, the system **100** may include a load balancer **110**, which may be configured to receive an I/O request **120** and to send the I/O request **120** to a web server **130**. The load balancer **110** may distribute I/O requests across multiple web servers **130** as part of a distributed storage system, which may be located in multiple different physical locations and/or may include multiple web servers **130** in a single physical location. Similarly, the web server **130** may be in communication with one or more storage systems **140**. The storage systems **140** may be configured to store data in PMR or SMR configurations. For example, a storage system **140** may include a storage server **142** in communication with one or more PMR drives **144**. Additionally, the web server **130** may be in communication with a storage system including an SMR storage server **146** that is in communication with one or more SMR drives **148**, as described in more detail in reference to FIG. 2. The system **100** may be configured, for example, through execution of software by the SMR storage server **146**, to execute the I/O request **120** on data stored in the SMR drives **148**.

(20) The system **100** may map data stored in the storage systems **140** to volume data by making

reference to a volume service **150**. In this way, the I/O request **120** may include a reference to a volume identifier, which may be mapped to a location in the storage of the storage system **140**, in a volume metadata database **160**. In an example, a read request may reference a volume identifier, which may be mapped to a zone segment, a segment pointer, and a read size, in an entry in the volume metadata database **160**.

(21) In this way, the system **100** may manage SMR drives **148** as part of a distributed data storage system. Advantageously, the system **100** may permit a back end data storage subsystem to benefit from the increased storage capacity of the SMR drives **148**, while maintaining integration with a front end block volume system. As described in more detail in reference to FIGS. **3-6**, the database system **100** may receive I/O requests **120**, for example, as part of data replication or backup, and may implement the I/O requests **120** by reading, writing, or defragmenting PMR drives **144** and/or SMR drives **148**.

(22) FIG. **2** illustrates an example shingled magnetic recording SMR drive **200**, in accordance with one or more embodiments. The physical processes and data structures of the SMR drive **200** are based on overwriting magnetic tracks **210** on disks of the SMR drive **200**. In a PMR drive, magnetic tracks are not overwritten. As such, data stored in the PMR drive may be accessed randomly (e.g., a PMR drive may be used as random-access hard drive storage). By contrast, SMR drives, being able to write by appending data rather than random-access, are typically configured for reading data efficiently that are unlikely to be modified.

(23) Data in the SMR drive **200** may be stored using a block system, rather than a file system, as part of the append-only I/O system. As such, data blocks may be written to zones **220** of a fixed size (e.g., 256 MB), which may be organized into zone segments **230**. Physically, each zone may be described by a number of magnetic tracks **210** having a fixed write width **240**, that are partially overwritten to provide a data track **250**. The data track may be described by a data track pitch **260**, which may be a parameter of the SMR drive **200** that depends, for example, on specifications of a read/write head **270** of the SMR drive **200**. With a narrower data track pitch **260**, the overall capacity of the SMR drive **200** increases, limited by factors including the ability of the read/write head **270** to accurately read the data track **250**.

(24) The magnetic tracks **210** are illustrated having a fixed width in order to demonstrate the hierarchical organization of the zones **220** and the zone segments **230**. Rather than indexing data to a single position in a discrete track, as is done in PMR drives, the SMR drive **200** may organize data in the zones **220** extending across multiple magnetic tracks **210**, indexed by pointers indicating the starting point of the zone segment **230**, the starting point of the zone **220**, and/or a location within the zone **220** where the data is to be found. For example, a pointer may include a write pointer, describing the last-written position in the zone **220**, where subsequent write operations will append new data.

(25) In some cases, the zone **220** may be described by one or more statuses that may govern whether data may be written to the zone **220**. For example, the statuses of the zone **220** may include, but are not limited to, “empty,” “full,” “open,” “closed,” or “finished.” An empty status may indicate that the zone **220** is able to be written from the beginning of the magnetic tracks **210** of the zone **220** (e.g., the write pointer for the zone **220** is at the start of the zone **220**). Similarly, a full status may indicate that the zone **220** is not available for additional write operations (e.g., the write pointer for the zone **220** is at the end or near the end of the zone **220**). The open status may indicate that the zone **220** is available for write operations (e.g., the zone **220** has capacity available for writing data to the magnetic tracks **210** of the zone **220** and the SMR drive **220** has resources assigned to implement write operations). Closed status, by contrast, may indicate that the zone **220** is not available for write operations, but may be opened. Finished status may indicate that data will not be written to the zone **220** (e.g., the write pointer for the zone **220** may be moved to the end of the zone to prevent further write operations). A finished zone may be opened by setting the zone **220** to an empty status (e.g., by moving the write pointer to the start of the zone **220**).

(26) Organizing the zones **220** into the zone segment **230** may permit the zones **220** to be managed together in such a way that a segment pointer may be defined, describing a write position for the zone segment **230**, rather than a write position for each zone. Within the zones **220** of the zone segment **230**, zone-specific write positions may follow the statuses described above, while the segment pointer may describe a write position in the first available open zone of the zone segment **230**. For example, the zones **220** may be organized in the zone segment **230** sequentially. In this way, data may be written in blocks across multiple zones **220** in the zone segment **230**. Organizing zones in a sequential manner may permit the SMR drive **200** to better accommodate write requests referencing data the length of which exceeds the capacity of the open zone. The SMR drive **200** may open multiple zones in a zone segment, and may define multiple write operations based on a single write request. As described in more detail in reference to FIG. 3, below, write operations may be implemented by identifying a segment pointer indicating a location to start writing data in a zone **220** of the zone segment **230**.

(27) FIG. 3 illustrates an example flow **300** for writing object data to a zone segment, in accordance with one or more embodiments. The operations of the flow can be implemented as hardware circuitry and/or stored as computer-readable instructions on a non-transitory computer-readable medium of a computer system, such as the system **100** of FIG. 1. As implemented, the instructions represent modules that include circuitry or code executable by a processor(s) of the computer system. The execution of such instructions configures the computer system to perform the specific operations described herein. Each circuitry or code in combination with the processor performs the respective operation(s). While the operations are illustrated in a particular order, it should be understood that no particular order is necessary and that one or more operations may be omitted, skipped, and/or reordered.

(28) In an example, the flow **300** includes an operation **310**, where the computer system receives a write request. The write request (e.g., I/O request **120** of FIG. 1) may be received from distributed data storage system managing data replication, archive, or other storage, based, for example, on a user request to store data in object storage, rather than in a block volume system. As part of receiving the write request, the computer system may forward the request to other networked systems (e.g., as part of an IaaS infrastructure system), as described in more detail in reference to FIG. 1. For example, the write request may be forwarded to a web server (e.g., web server **130** of FIG. 1), which may select a storage server (e.g., SMR storage server **146** of FIG. 1) to handle the write request.

(29) In an example, the flow **300** optionally includes an operation **320**, where the computer system scans the zone segment for empty or open zones. The write request may include a data chunk of a defined length, where the chunk is a block of data to be written to the zone segment (e.g., zone segment **230** of FIG. 2). As described in more detail in reference to FIG. 2, the zone segment may include multiple zones (e.g., zone **220** of FIG. 2) providing a storage capacity equivalent to the integer multiple of the capacity of the zones. One or more zones in the zone segment may be empty or open, and may therefore be available for write operations. The zones may be organized in the zone segment in sequence, such that an open zone may be followed by an empty zone. It should be understood that an empty zone does not necessarily indicate a zone for which the magnetic tracks are blank, but rather may indicate a zone for which the write position for the zone is at the start of the zone.

(30) In an example, the flow **300** optionally includes an operation **330**, where the computer system allocates a new zone segment to the volume. In some cases, the scan of the zone segment may not identify an empty or open zone before the terminal zone of the zone segment. In some embodiments, the system may allocate a new zone segment to the volume referenced in the write request (e.g., by a volume identifier) to increase the capacity available to write data for the referenced volume. The new zone segment may be allocated from a free pool of zone segments maintained for the purpose of reducing the proportion of write requests that are rejected for lack of

capacity. In some embodiments, the new zone segment may be allocated from the same SMR drive storing the zone segment, or may be allocated from a different SMR drive, coordinated through an SMR storage server (e.g., SMR storage server **146** of FIG. **1**).

(31) In an example, the flow **300** includes an operation **340**, where the computer system identifies a segment pointer. Within the zone segment (or the new zone segment) the system may identify the segment pointer by locating the first zone of the zone segment described by an open or empty status. For example, where the zones are organized sequentially in the zone segment, the segment pointer may correspond to the write position of the first open zone. In another example, the segment pointer may correspond to the start position of the first empty zone, where the zone segment does not include an open zone.

(32) In an example, the flow **300** includes an operation **350**, where the computer system writes data to a zone of the zone segments according to the segment pointer. Starting from the segment pointer, the system may write data referenced by the write request to the zone corresponding to the segment pointer. The write operation may include an initial check by the system whether the zone corresponding to the segment pointer has capacity for the entire chunk referenced by the write request. In cases where the zone has insufficient capacity, the system may define multiple blocks, such that the chunk may be split across multiple zones of the zone segment. In an illustrative example, the zones of the zone segment may be organized sequentially, and a chunk may be split into a first block and a second block in such a way that the size of the first block corresponds to the write capacity of the zone corresponding to the segment pointer. The first block may be written starting at the segment pointer and the second block may be written to the next zone. While this example describes zones in sequence, method **300** may be implemented with other approaches.

(33) In an example, the flow **300** includes an operation **360**, where the computer system updates the segment pointer. Following completion of the write operation, which may include multiple write operations where the data is written to multiple zones, the system may define the segment pointer to the end position of the data. In this way, the updated segment pointer may correspond to a write position in a different zone of the zone segment, relative to the zone corresponding to the segment pointer preceding the receipt of the write request.

(34) In an example, the flow **300** includes an operation **370**, where the computer system updates the volume metadata. As described in more detail in reference to FIG. **1**, the data written to the zone segment may be mapped to a block volume through volume metadata. The volume metadata may include the segment pointer, and, as such, updating the volume metadata may include updating the segment pointer in a volume metadata database (e.g., volume metadata database **160** of FIG. **1**), as well as updating the volume-zone segment mapping in cases where new zone segments have been allocated, as part of operations **320** and **330**. In some examples, the volume metadata may store the start location pointer from which the data was written which corresponds to the segment pointer value before it was updated in operation **360**.

(35) In an example, the flow **300** includes an operation **380**, where the computer system outputs a status of the write request. In some embodiments, the system may be configured to output various different types of status reports. For example, resource metadata, such as usage and capacity data may be collected and managed as part of improving IaaS system operations. As another example, the success or failure of the write request may be returned to the system, such that the system may repeat some or all of the operations of the method **300**. As another example, the system may output a status to a user of an IaaS console interface, where the write request originated with a human user.

(36) Zone segment organization may provide a significant technical improvement for the operation of a database system that incorporates SMR drives. By implementing the techniques of writing to a zone segment, including multiple zones, data may be written across zones, and capacity of multiple zones may be shared. In this way, data replication, storage, and backup operations, facilitated by volume-metadata mapped to zone segment pointers, may improve the efficiency of SMR write and re/write operations.

(37) FIG. 4 illustrates an example flow **400** for reading data from a zone segment, in accordance with one or more embodiments. The operations of the flow can be implemented as hardware circuitry and/or stored as computer-readable instructions on a non-transitory computer-readable medium of a computer system, such as the system **100** of FIG. 1. As implemented, the instructions represent modules that include circuitry or code executable by a processor(s) of the computer system. The execution of such instructions configures the computer system to perform the specific operations described herein. Each circuitry or code in combination with the processor performs the respective operation(s). While the operations are illustrated in a particular order, it should be understood that no particular order is necessary and that one or more operations may be omitted, skipped, and/or reordered.

(38) In an example, the flow **400** includes an operation **410**, where the computer system receives a read request. As described in more detail in reference to FIG. 1, a read request may reference a volume and a data block to be read. The volume may be referenced by a volume identifier and a block identifier, which may be mapped in volume metadata to an SMR drive. SMR drives may define zone segments including zones, which may be organized sequentially in the zone segments. In this way, the system may receive the read request as part of a data access operation, for example, as part of a data restore process, and may be configured to access multiple zones of a zone segment to read the data.

(39) In an example, the flow **400** includes an operation **420**, where the computer system identifies a stored pointer that was stored in the volume metadata (e.g., when the volume metadata is updated at **370** of FIG. 3) and a read length. The stored pointer and the read length may correspond to a volume identified in the read request. The volume identifier may be mapped to a zone segment, and to a zone of the zone segment, by the stored pointer. In contrast to the method described in reference to FIG. 3, the stored pointer of the flow **400** may reference a read position, rather than a write position. Even so, it similarly may be understood to reference to a start position in the zone segment for the SMR drive to read data.

(40) The read request may also include reference to data to be read, which may be mapped to a read length on the SMR drive. The read length may correspond to a dimension or length of track (e.g., data track **250** of FIG. 2) to be read, such that the system may determine whether the read length overlaps multiple zones in a zone segment. For example, where zones are organized sequentially in a zone segment, the system may indicate a first read position at the segment pointer and a second read position at the start position of the next zone in the zone segment.

(41) In an example, the flow **400** optionally includes an operation **430**, where the computer system determines single read commands for each zone in the zone segment. Where the read length, starting from the segment pointer identified in operation **420**, extends over multiple zones, the system may determine multiple read commands for different zones of the zone segment. The multiple single read commands may be cross-referenced, such that the data returned by the read command may be reconstructed according to the read length determined in operation **420**. In this way, data written across multiple zones may be read by a single read request referencing a zone segment, rather than multiple read requests addressed at individual zones.

(42) In an example, the flow **400** includes an operation **440**, where the computer system issues a read command to the SMR drive. As described in more detail in reference to FIG. 1, the system may manage data storage via storage servers (e.g., storage server **142** or SMR server **146** of FIG. 1). The read command may include the segment pointer and read length determined in operation **420**, and the system may issue the read command to an SMR storage server referenced by mapping metadata, such as the volume metadata as described in more detail in reference to FIG. 1.

(43) In an example, the flow **400** includes an operation **450**, where the computer system outputs read data and a status of the read request. Output of the data may include communicating the data from the SMR drive to a destination referenced in the read request. For example, the system may output the data to a volume storage system, or as part of data replication operations from one object

storage drive to another object storage drive. Similarly to the flow **300**, the status may include success/failure information for the request, as well as other information including, but not limited to, metadata describing the data and zones referenced in the read command.

(44) Facilitating read commands using the operations of the flow **400** may provide a significant technical improvement for the operation of a database system that incorporates SMR drives. By implementing the techniques of serving read requests addressed at multiple zones organized into zone segments, data may be more efficiently read from SMR drives without multiple read requests generated for each zone in the zone segment.

(45) FIG. 5 illustrates an example flow **500** for deleting data from a zone segment, in accordance with one or more embodiments. The operations of the flow can be implemented as hardware circuitry and/or stored as computer-readable instructions on a non-transitory computer-readable medium of a computer system, such as the system **100** of FIG. 1. As implemented, the instructions represent modules that include circuitry or code executable by a processor(s) of the computer system. The execution of such instructions configures the computer system to perform the specific operations described herein. Each circuitry or code in combination with the processor performs the respective operation(s). While the operations are illustrated in a particular order, it should be understood that no particular order is necessary and that one or more operations may be omitted, skipped, and/or reordered.

(46) In an example, the flow **500** includes an operation **510**, where the computer system receives a delete request. In contrast to read or write operations, delete operations may proceed without reference to a segment pointer as a start position. Instead, the delete request may include volume identifier information, which may be mapped to object storage in one or more drives, such as SMR drives as described in more detail in reference to FIG. 2 (e.g., SMR drive **200** of FIG. 2).

(47) In an example, the flow **500** includes an operation **520**, where the computer system reads mapping metadata for the volume object. As described in more detail in reference to FIG. 1, mapping metadata may reference the zone segment storing data for the block volume subject to the delete request. In this way, the system may identify the zone segment or zone segments allocated to the volume identified in the delete request by making reference to the mapping metadata.

(48) In an example, the flow **500** includes an operation **530**, where the computer system updates a state of the volume object. Updating the state of the volume object may include, but is not limited to, sending an instruction to a storage server (e.g., SMR storage server **146** of FIG. 1) to modify the status of the zones included in an allocated zone segment from “full” or “open” to “empty,” such that for subsequent write requests, the segment pointer corresponds to a start position of the first zone in the zone segment. Additionally or alternatively, the system may modify volume metadata in a volume metadata database (e.g., volume metadata database **160** of FIG. 1) to remove mapping information, such that a volume identifier included in subsequent I/O requests (e.g., I/O request **120** of FIG. 1) is not associated with zone segments in an SMR drive. For example, in the operations of the flow **300** described in reference to FIG. 3, identifying the segment pointer may return that no zone segments are allocated to the referenced volume identifier. In this way, delete requests may be addressed to a portion of a zone of a zone segment. Rather than erasing an entire zone, which may permit the SMR drive to operate in a manner more similar to a random-access drive, rather than an append-only drive, except that write requests may not write to partially overwritten tracks.

(49) In an example, the flow **500** includes an operation **540**, where the computer system outputs a status of the delete request. As with the preceding flows of FIGS. 3-4, the system may return a status of the delete request to indicate that the resources addressed by the delete request have been released, for example, to a central metadata management system. Additionally or alternatively, the system may return a status success/failure message to the originator of the delete request, to indicate whether the request should be resent or will be re-attempted.

(50) FIG. 6 illustrates an example flow **600** for defragmenting a zone segment, in accordance with

one or more embodiments. The operations of the flow can be implemented as hardware circuitry and/or stored as computer-readable instructions on a non-transitory computer-readable medium of a computer system, such as the system **100** of FIG. **1**. As implemented, the instructions represent modules that include circuitry or code executable by a processor(s) of the computer system. The execution of such instructions configures the computer system to perform the specific operations described herein. Each circuitry or code in combination with the processor performs the respective operation(s). While the operations are illustrated in a particular order, it should be understood that no particular order is necessary and that one or more operations may be omitted, skipped, and/or reordered.

(51) In addition to improving read/write operations on SMR drives, in some embodiments, defragmentation operations may also be improved through organizing zones of an SMR drive into zone segments, and remapping data stored in zones to reduce inefficient use of storage.

Advantageously, defragmenting zone segments may permit SMR drives to function more like random-access drives, rather than strictly as append-only drives, but will still be limited to append-only write operations. In an illustrative example, delete requests, as described in more detail in reference to FIG. **5**, may modify mapping metadata rather than data written to drives directly. Over the course of operation, an SMR drive (e.g., SMR drive **200** of FIG. **2**) may use a fraction of the capacity of the zone segment, as when data has been deleted that begins in one zone and ends in a subsequent zone of the zone segment. As data in partially overlapped tracks cannot be rewritten in SMR drives, the unused capacity remaining may be unusable for write operations, which may introduce inefficiency and under-utilization of storage capacity on the SMR drive. Defragmenting the SMR drive may permit the system to release unused storage capacity, by rewriting data into new zones and remapping volume metadata.

(52) In an example, the flow **600** includes an operation **610**, where the computer system ascertains a usage fraction of a zone segment. As part of defragmentation, the system may implement automatic (e.g., without human intervention) processes based on a usage fraction exceeding a threshold to trigger the defragment operations (e.g., periodic checking of the usage fraction, refreshing the usage fraction subsequent to a delete request, or the like). To that end, the system may access the volume metadata, to ascertain the usage fraction. The usage fraction may describe the portion of storage capacity in the zones of the zone segment that are mapped in volume metadata. For example, volume metadata may reference starting points, ending points, chunk lengths, or the like, for data stored in the zone segment, amounting to only a portion of the total capacity of the zone segment, based, for example, on a zone segment having a fixed capacity.

(53) In an example, the flow **600** includes an operation **620**, where the computer system compares the usage fraction to a usage threshold. The usage threshold may be expressed in terms of a usage fraction. For example, where the system is configured to implement defragmentation when the usage fraction drops below 50%, the usage threshold may be expressed as 0.5.

(54) In an example, the flow **600** includes an operation **630**, where the computer system identifies data stored in the zones of the zone segment. When the defragmentation process has been triggered as a result of the comparison of operation **620**, the system may identify, using volume metadata and/or by scanning the zone segment, the data that is to be moved during defragmentation of the zone segment. While defragmentation of a random-access drive may permit moving some data while leaving other data in place on the drive, append-only writing of SMR drives may include identifying and remapping data to new zones, and erasing the zones from which the data is copied.

(55) In an example, the flow **600** includes an operation **640**, where the computer system maps stored data to empty or open zones. In some instances, this is a two-part operation, where the first part occurs here at operation **640** (e.g., this is similar to how the write destination location is determined at operation **340** of FIG. **3**) and the second part happens after operation **650** (described later). Mapping data that was identified in operation **630** may include modifying volume metadata, such that the data associated with a volume is mapped to new zones of the zone segment, to newly

allocated zones, or to buffer storage, which may be in a zone segment maintained free for the use of defragmentation on the SMR drive. Similarly, random access memory may be used for the buffer memory. Remapping may include determining a configuration of the stored data that for subsequent writing of the stored data to a zone segment.

(56) In an example, the flow **600** includes an operation **650**, where the computer system writes stored data to the empty or open zones. Writing the stored data may include the operations of the flow **300**, as described in more detail in reference to FIG. 3. For example, the system may send a write command to an SMR storage server (e.g., SMR storage server **146** of FIG. 1), to cause the stored data to be written to one or more empty or open zones. Where the write operation is performed on zones that are not reserved for defragmentation, the data may be left in where written, and the mapping may be maintained as provided by operation **640**.

(57) In an example, the second part of the two-part operation noted above at operation **640** occurs at operation **655** of flow **600**, after the data is copied at operation **650**. For example, once the data is safe in the new location, the volume metadata is updated (e.g., similar to operation **370** of FIG. 3).

(58) In an example, the flow **600** includes an operation **660**, where the computer system resets the source zone of the stored data. In addition, the system may reset the source zones of the zone segment for subsequent I/O operations (e.g., I/O operations **120** of FIG. 1). Resetting the source zones may include updating the status of the zone from “full” to “empty” or other modifications to reflect the new configuration of data in the zone segment.

(59) In an example, the flow **600** optionally includes an operation **670**, where the computer system releases excess zones. As described for operation **330** of the flow **300** in reference to FIG. 3, write operations for which the data to be written exceeds the capacity of the referenced zone segment may include allocating a new zone segment to provide additional storage capacity. In this way, defragmentation may make it such that the data may be stored in a single zone segment, by reducing the unused portion of data capacity in the zone segment. As such, zone segments may be released that were previously used to store data for the zone segment. In some embodiments, excess zones may be identified as part of mapping operations of **640**, and released following writing the stored data according to the new mapping as described in reference to operation **650**.

(60) In an example, the flow **600** optionally includes an operation **680**, where the computer system outputs the status of the defragment process. As with the output operations of the preceding FIGS. 3-5, operation **680** may include reporting resource usage status to a central database of resource metadata, and/or indicating a success/failure of the defragmentation process. In some embodiments, operation **680** may also include ascertaining a new usage fraction of the zone or resetting a timer or other trigger condition.

(61) Defragmenting on a zone-segment basis may provide a significant technical improvement for the operation of a database system that incorporates SMR drives. By implementing the techniques of zone-segment defragmentation, efficient usage of SMR drive capacity may be improved, at least by improving the usage fraction of zone segments, where data may be written across zones of the zone segment. In this way, limitations introduced by append-only write operations, which are characteristic of SMR drives, may be reduced.

(62) As noted above, infrastructure as a service (IaaS) is one particular type of cloud computing. IaaS can be configured to provide virtualized computing resources over a public network (e.g., the Internet). In an IaaS model, a cloud computing provider can host the infrastructure components (e.g., servers, storage devices, network nodes (e.g., hardware), deployment software, platform virtualization (e.g., a hypervisor layer), or the like). In some cases, an IaaS provider may also supply a variety of services to accompany those infrastructure components (e.g., billing, monitoring, logging, security, load balancing and clustering, etc.). Thus, as these services may be policy-driven, IaaS users may be able to implement policies to drive load balancing to maintain application availability and performance.

(63) In some instances, IaaS customers may access resources and services through a wide area

network (WAN), such as the Internet, and can use the cloud provider's services to install the remaining elements of an application stack. For example, the user can log in to the IaaS platform to create virtual machines (VMs), install operating systems (OSs) on each VM, deploy middleware such as databases, create storage buckets for workloads and backups, and even install enterprise software into that VM. Customers can then use the provider's services to perform various functions, including balancing network traffic, troubleshooting application issues, monitoring performance, managing disaster recovery, etc.

(64) In most cases, a cloud computing model will require the participation of a cloud provider. The cloud provider may, but need not be, a third-party service that specializes in providing (e.g., offering, renting, selling) IaaS. An entity might also opt to deploy a private cloud, becoming its own provider of infrastructure services.

(65) In some examples, IaaS deployment is the process of putting a new application, or a new version of an application, onto a prepared application server or the like. It may also include the process of preparing the server (e.g., installing libraries, daemons, etc.). This is often managed by the cloud provider, below the hypervisor layer (e.g., the servers, storage, network hardware, and virtualization). Thus, the customer may be responsible for handling (OS), middleware, and/or application deployment (e.g., on self-service virtual machines (e.g., that can be spun up on demand) or the like.

(66) In some examples, IaaS provisioning may refer to acquiring computers or virtual hosts for use, and even installing needed libraries or services on them. In most cases, deployment does not include provisioning, and the provisioning may need to be performed first.

(67) In some cases, there are two different problems for IaaS provisioning. First, there is the initial challenge of provisioning the initial set of infrastructure before anything is running. Second, there is the challenge of evolving the existing infrastructure (e.g., adding new services, changing services, removing services, etc.) once everything has been provisioned. In some cases, these two challenges may be addressed by enabling the configuration of the infrastructure to be defined declaratively. In other words, the infrastructure (e.g., what components are needed and how they interact) can be defined by one or more configuration files. Thus, the overall topology of the infrastructure (e.g., what resources depend on which, and how they each work together) can be described declaratively. In some instances, once the topology is defined, a workflow can be generated that creates and/or manages the different components described in the configuration files.

(68) In some examples, an infrastructure may have many interconnected elements. For example, there may be one or more virtual private clouds (VPCs) (e.g., a potentially on-demand pool of configurable and/or shared computing resources), also known as a core network. In some examples, there may also be one or more security group rules provisioned to define how the security of the network will be set up and one or more virtual machines (VMs). Other infrastructure elements may also be provisioned, such as a load balancer, a database, or the like. As more and more infrastructure elements are desired and/or added, the infrastructure may incrementally evolve.

(69) In some instances, continuous deployment techniques may be employed to enable deployment of infrastructure code across various virtual computing environments. Additionally, the described techniques can enable infrastructure management within these environments. In some examples, service teams can write code that is desired to be deployed to one or more, but often many, different production environments (e.g., across various different geographic locations, sometimes spanning the entire world). However, in some examples, the infrastructure on which the code will be deployed must first be set up. In some instances, the provisioning can be done manually, a provisioning tool may be utilized to provision the resources, and/or deployment tools may be utilized to deploy the code once the infrastructure is provisioned.

(70) FIG. 7 is a block diagram **700** illustrating an example pattern of an IaaS architecture, according to at least one embodiment. Service operators **702** can be communicatively coupled to a secure host tenancy **704** that can include a virtual cloud network (VCN) **706** and a secure host

subnet **708**. In some examples, the service operators **702** may be using one or more client computing devices, which may be portable handheld devices (e.g., an iPhone®, cellular telephone, an iPad®, computing tablet, a personal digital assistant (PDA)) or wearable devices (e.g., a Google Glass® head mounted display), running software such as Microsoft Windows Mobile®, and/or a variety of mobile operating systems such as iOS, Windows Phone, Android, BlackBerry 7, Palm OS, and the like, and being Internet, e-mail, short message service (SMS), Blackberry®, or other communication protocol enabled. Alternatively, the client computing devices can be general purpose personal computers including, by way of example, personal computers and/or laptop computers running various versions of Microsoft Windows®, Apple Macintosh®, and/or Linux operating systems. The client computing devices can be workstation computers running any of a variety of commercially-available UNIX® or UNIX-like operating systems, including without limitation the variety of GNU/Linux operating systems, such as for example, Google Chrome OS. Alternatively, or in addition, client computing devices may be any other electronic device, such as a thin-client computer, an Internet-enabled gaming system (e.g., a Microsoft Xbox gaming console with or without a Kinect® gesture input device), and/or a personal messaging device, capable of communicating over a network that can access the VCN **706** and/or the Internet.

(71) The VCN **706** can include a local peering gateway (LPG) **710** that can be communicatively coupled to a secure shell (SSH) VCN **712** via an LPG **710** contained in the SSH VCN **712**. The SSH VCN **712** can include an SSH subnet **714**, and the SSH VCN **712** can be communicatively coupled to a control plane VCN **716** via the LPG **710** contained in the control plane VCN **716**. Also, the SSH VCN **712** can be communicatively coupled to a data plane VCN **718** via an LPG **710**. The control plane VCN **716** and the data plane VCN **718** can be contained in a service tenancy **719** that can be owned and/or operated by the IaaS provider.

(72) The control plane VCN **716** can include a control plane demilitarized zone (DMZ) tier **720** that acts as a perimeter network (e.g., portions of a corporate network between the corporate intranet and external networks). The DMZ-based servers may have restricted responsibilities and help keep security breaches contained. Additionally, the DMZ tier **720** can include one or more load balancer (LB) subnet(s) **722**, a control plane app tier **724** that can include app subnet(s) **726**, a control plane data tier **728** that can include database (DB) subnet(s) **730** (e.g., frontend DB subnet(s) and/or backend DB subnet(s)). The LB subnet(s) **722** contained in the control plane DMZ tier **720** can be communicatively coupled to the app subnet(s) **726** contained in the control plane app tier **724** and an Internet gateway **734** that can be contained in the control plane VCN **716**, and the app subnet(s) **726** can be communicatively coupled to the DB subnet(s) **730** contained in the control plane data tier **728** and a service gateway **736** and a network address translation (NAT) gateway **738**. The control plane VCN **716** can include the service gateway **736** and the NAT gateway **738**.

(73) The control plane VCN **716** can include a data plane mirror app tier **740** that can include app subnet(s) **726**. The app subnet(s) **726** contained in the data plane mirror app tier **740** can include a virtual network interface controller (VNIC) **742** that can execute a compute instance **744**. The compute instance **744** can communicatively couple the app subnet(s) **726** of the data plane mirror app tier **740** to app subnet(s) **726** that can be contained in a data plane app tier **746**.

(74) The data plane VCN **718** can include the data plane app tier **746**, a data plane DMZ tier **748**, and a data plane data tier **750**. The data plane DMZ tier **748** can include LB subnet(s) **722** that can be communicatively coupled to the app subnet(s) **726** of the data plane app tier **746** and the Internet gateway **734** of the data plane VCN **718**. The app subnet(s) **726** can be communicatively coupled to the service gateway **736** of the data plane VCN **718** and the NAT gateway **738** of the data plane VCN **718**. The data plane data tier **750** can also include the DB subnet(s) **730** that can be communicatively coupled to the app subnet(s) **726** of the data plane app tier **746**.

(75) The Internet gateway **734** of the control plane VCN **716** and of the data plane VCN **718** can be communicatively coupled to a metadata management service **752** that can be communicatively

coupled to public Internet **754**. Public Internet **754** can be communicatively coupled to the NAT gateway **738** of the control plane VCN **716** and of the data plane VCN **718**. The service gateway **736** of the control plane VCN **716** and of the data plane VCN **718** can be communicatively couple to cloud services **756**.

(76) In some examples, the service gateway **736** of the control plane VCN **716** or of the data plane VCN **718** can make application programming interface (API) calls to cloud services **756** without going through public Internet **754**. The API calls to cloud services **756** from the service gateway **736** can be one-way: the service gateway **736** can make API calls to cloud services **756**, and cloud services **756** can send requested data to the service gateway **736**. But, cloud services **756** may not initiate API calls to the service gateway **736**.

(77) In some examples, the secure host tenancy **704** can be directly connected to the service tenancy **719**, which may be otherwise isolated. The secure host subnet **708** can communicate with the SSH subnet **714** through an LPG **710** that may enable two-way communication over an otherwise isolated system. Connecting the secure host subnet **708** to the SSH subnet **714** may give the secure host subnet **708** access to other entities within the service tenancy **719**.

(78) The control plane VCN **716** may allow users of the service tenancy **719** to set up or otherwise provision desired resources. Desired resources provisioned in the control plane VCN **716** may be deployed or otherwise used in the data plane VCN **718**. In some examples, the control plane VCN **716** can be isolated from the data plane VCN **718**, and the data plane mirror app tier **740** of the control plane VCN **716** can communicate with the data plane app tier **746** of the data plane VCN **718** via VNICs **742** that can be contained in the data plane mirror app tier **740** and the data plane app tier **746**.

(79) In some examples, users of the system, or customers, can make requests, for example create, read, update, or delete (CRUD) operations, through public Internet **754** that can communicate the requests to the metadata management service **752**. The metadata management service **752** can communicate the request to the control plane VCN **716** through the Internet gateway **734**. The request can be received by the LB subnet(s) **722** contained in the control plane DMZ tier **720**. The LB subnet(s) **722** may determine that the request is valid, and in response to this determination, the LB subnet(s) **722** can transmit the request to app subnet(s) **726** contained in the control plane app tier **724**. If the request is validated and requires a call to public Internet **754**, the call to public Internet **754** may be transmitted to the NAT gateway **738** that can make the call to public Internet **754**. Memory that may be desired to be stored by the request can be stored in the DB subnet(s) **730**.

(80) In some examples, the data plane mirror app tier **740** can facilitate direct communication between the control plane VCN **716** and the data plane VCN **718**. For example, changes, updates, or other suitable modifications to configuration may be desired to be applied to the resources contained in the data plane VCN **718**. Via a VNIC **742**, the control plane VCN **716** can directly communicate with, and can thereby execute the changes, updates, or other suitable modifications to configuration to, resources contained in the data plane VCN **718**.

(81) In some embodiments, the control plane VCN **716** and the data plane VCN **718** can be contained in the service tenancy **719**. In this case, the user, or the customer, of the system may not own or operate either the control plane VCN **716** or the data plane VCN **718**. Instead, the IaaS provider may own or operate the control plane VCN **716** and the data plane VCN **718**, both of which may be contained in the service tenancy **719**. This embodiment can enable isolation of networks that may prevent users or customers from interacting with other users', or other customers' resources. Also, this embodiment may allow users or customers of the system to store databases privately without needing to rely on public Internet **754**, which may not have a desired level of security, for storage.

(82) In other embodiments, the LB subnet(s) **722** contained in the control plane VCN **716** can be configured to receive a signal from the service gateway **736**. In this embodiment, the control plane VCN **716** and the data plane VCN **718** may be configured to be called by a customer of the IaaS

provider without calling public Internet **754**. Customers of the IaaS provider may desire this embodiment since database(s) that the customers use may be controlled by the IaaS provider and may be stored on the service tenancy **719**, which may be isolated from public Internet **754**.

(83) FIG. **8** is a block diagram **800** illustrating another example pattern of an IaaS architecture, according to at least one embodiment. Service operators **802** (e.g. service operators **702** of FIG. **7**) can be communicatively coupled to a secure host tenancy **804** (e.g. the secure host tenancy **704** of FIG. **7**) that can include a virtual cloud network (VCN) **806** (e.g. the VCN **706** of FIG. **7**) and a secure host subnet **808** (e.g. the secure host subnet **708** of FIG. **7**). The VCN **806** can include a local peering gateway (LPG) **810** (e.g. the LPG **710** of FIG. **7**) that can be communicatively coupled to a secure shell (SSH) VCN **812** (e.g. the SSH VCN **712** of FIG. **7**) via an LPG **710** contained in the SSH VCN **812**. The SSH VCN **812** can include an SSH subnet **814** (e.g. the SSH subnet **714** of FIG. **7**), and the SSH VCN **812** can be communicatively coupled to a control plane VCN **816** (e.g. the control plane VCN **716** of FIG. **7**) via an LPG **810** contained in the control plane VCN **816**. The control plane VCN **816** can be contained in a service tenancy **819** (e.g. the service tenancy **719** of FIG. **7**), and the data plane VCN **818** (e.g. the data plane VCN **718** of FIG. **7**) can be contained in a customer tenancy **821** that may be owned or operated by users, or customers, of the system.

(84) The control plane VCN **816** can include a control plane DMZ tier **820** (e.g. the control plane DMZ tier **720** of FIG. **7**) that can include LB subnet(s) **822** (e.g. LB subnet(s) **722** of FIG. **7**), a control plane app tier **824** (e.g. the control plane app tier **724** of FIG. **7**) that can include app subnet(s) **826** (e.g. app subnet(s) **726** of FIG. **7**), a control plane data tier **828** (e.g. the control plane data tier **728** of FIG. **7**) that can include database (DB) subnet(s) **830** (e.g. similar to DB subnet(s) **730** of FIG. **7**). The LB subnet(s) **822** contained in the control plane DMZ tier **820** can be communicatively coupled to the app subnet(s) **826** contained in the control plane app tier **824** and an Internet gateway **834** (e.g. the Internet gateway **734** of FIG. **7**) that can be contained in the control plane VCN **816**, and the app subnet(s) **826** can be communicatively coupled to the DB subnet(s) **830** contained in the control plane data tier **828** and a service gateway **836** (e.g. the service gateway of FIG. **7**) and a network address translation (NAT) gateway **838** (e.g. the NAT gateway **738** of FIG. **7**). The control plane VCN **816** can include the service gateway **836** and the NAT gateway **838**.

(85) The control plane VCN **816** can include a data plane mirror app tier **840** (e.g. the data plane mirror app tier **740** of FIG. **7**) that can include app subnet(s) **826**. The app subnet(s) **826** contained in the data plane mirror app tier **840** can include a virtual network interface controller (VNIC) **842** (e.g. the VNIC of **742**) that can execute a compute instance **844** (e.g. similar to the compute instance **744** of FIG. **7**). The compute instance **844** can facilitate communication between the app subnet(s) **826** of the data plane mirror app tier **840** and the app subnet(s) **826** that can be contained in a data plane app tier **846** (e.g. the data plane app tier **746** of FIG. **7**) via the VNIC **842** contained in the data plane mirror app tier **840** and the VNIC **842** contained in the data plane app tier **846**.

(86) The Internet gateway **834** contained in the control plane VCN **816** can be communicatively coupled to a metadata management service **852** (e.g. the metadata management service **752** of FIG. **7**) that can be communicatively coupled to public Internet **854** (e.g. public Internet **754** of FIG. **7**). Public Internet **854** can be communicatively coupled to the NAT gateway **838** contained in the control plane VCN **816**. The service gateway **836** contained in the control plane VCN **816** can be communicatively couple to cloud services **856** (e.g. cloud services **756** of FIG. **7**).

(87) In some examples, the data plane VCN **818** can be contained in the customer tenancy **821**. In this case, the IaaS provider may provide the control plane VCN **816** for each customer, and the IaaS provider may, for each customer, set up a unique compute instance **844** that is contained in the service tenancy **819**. Each compute instance **844** may allow communication between the control plane VCN **816**, contained in the service tenancy **819**, and the data plane VCN **818** that is contained in the customer tenancy **821**. The compute instance **844** may allow resources, that are

provisioned in the control plane VCN **816** that is contained in the service tenancy **819**, to be deployed or otherwise used in the data plane VCN **818** that is contained in the customer tenancy **821**.

(88) In other examples, the customer of the IaaS provider may have databases that live in the customer tenancy **821**. In this example, the control plane VCN **816** can include the data plane mirror app tier **840** that can include app subnet(s) **826**. The data plane mirror app tier **840** can reside in the data plane VCN **818**, but the data plane mirror app tier **840** may not live in the data plane VCN **818**. That is, the data plane mirror app tier **840** may have access to the customer tenancy **821**, but the data plane mirror app tier **840** may not exist in the data plane VCN **818** or be owned or operated by the customer of the IaaS provider. The data plane mirror app tier **840** may be configured to make calls to the data plane VCN **818** but may not be configured to make calls to any entity contained in the control plane VCN **816**. The customer may desire to deploy or otherwise use resources in the data plane VCN **818** that are provisioned in the control plane VCN **816**, and the data plane mirror app tier **840** can facilitate the desired deployment, or other usage of resources, of the customer.

(89) In some embodiments, the customer of the IaaS provider can apply filters to the data plane VCN **818**. In this embodiment, the customer can determine what the data plane VCN **818** can access, and the customer may restrict access to public Internet **854** from the data plane VCN **818**. The IaaS provider may not be able to apply filters or otherwise control access of the data plane VCN **818** to any outside networks or databases. Applying filters and controls by the customer onto the data plane VCN **818**, contained in the customer tenancy **821**, can help isolate the data plane VCN **818** from other customers and from public Internet **854**.

(90) In some embodiments, cloud services **856** can be called by the service gateway **836** to access services that may not exist on public Internet **854**, on the control plane VCN **816**, or on the data plane VCN **818**. The connection between cloud services **856** and the control plane VCN **816** or the data plane VCN **818** may not be live or continuous. Cloud services **856** may exist on a different network owned or operated by the IaaS provider. Cloud services **856** may be configured to receive calls from the service gateway **836** and may be configured to not receive calls from public Internet **854**. Some cloud services **856** may be isolated from other cloud services **856**, and the control plane VCN **816** may be isolated from cloud services **856** that may not be in the same region as the control plane VCN **816**. For example, the control plane VCN **816** may be located in “Region 1,” and cloud service “Deployment 6,” may be located in Region 1 and in “Region 2.” If a call to Deployment 6 is made by the service gateway **836** contained in the control plane VCN **816** located in Region 1, the call may be transmitted to Deployment 6 in Region 1. In this example, the control plane VCN **816**, or Deployment 6 in Region 1, may not be communicatively coupled to, or otherwise in communication with, Deployment 6 in Region 2.

(91) FIG. 9 is a block diagram **900** illustrating another example pattern of an IaaS architecture, according to at least one embodiment. Service operators **902** (e.g. service operators **702** of FIG. 7) can be communicatively coupled to a secure host tenancy **904** (e.g. the secure host tenancy **704** of FIG. 7) that can include a virtual cloud network (VCN) **906** (e.g. the VCN **706** of FIG. 7) and a secure host subnet **908** (e.g. the secure host subnet **708** of FIG. 7). The VCN **906** can include an LPG **910** (e.g. the LPG **710** of FIG. 7) that can be communicatively coupled to an SSH VCN **912** (e.g. the SSH VCN **712** of FIG. 7) via an LPG **910** contained in the SSH VCN **912**. The SSH VCN **912** can include an SSH subnet **914** (e.g. the SSH subnet **714** of FIG. 7), and the SSH VCN **912** can be communicatively coupled to a control plane VCN **916** (e.g. the control plane VCN **716** of FIG. 7) via an LPG **910** contained in the control plane VCN **916** and to a data plane VCN **918** (e.g. the data plane **718** of FIG. 7) via an LPG **910** contained in the data plane VCN **918**. The control plane VCN **916** and the data plane VCN **918** can be contained in a service tenancy **919** (e.g. the service tenancy **719** of FIG. 7).

(92) The control plane VCN **916** can include a control plane DMZ tier **920** (e.g. the control plane

DMZ tier **720** of FIG. 7) that can include load balancer (LB) subnet(s) **922** (e.g. LB subnet(s) **722** of FIG. 7), a control plane app tier **924** (e.g. the control plane app tier **724** of FIG. 7) that can include app subnet(s) **926** (e.g. similar to app subnet(s) **726** of FIG. 7), a control plane data tier **928** (e.g. the control plane data tier **728** of FIG. 7) that can include DB subnet(s) **930**. The LB subnet(s) **922** contained in the control plane DMZ tier **920** can be communicatively coupled to the app subnet(s) **926** contained in the control plane app tier **924** and to an Internet gateway **934** (e.g. the Internet gateway **734** of FIG. 7) that can be contained in the control plane VCN **916**, and the app subnet(s) **926** can be communicatively coupled to the DB subnet(s) **930** contained in the control plane data tier **928** and to a service gateway **936** (e.g. the service gateway of FIG. 7) and a network address translation (NAT) gateway **938** (e.g. the NAT gateway **738** of FIG. 7). The control plane VCN **916** can include the service gateway **936** and the NAT gateway **938**.

(93) The data plane VCN **918** can include a data plane app tier **946** (e.g. the data plane app tier **746** of FIG. 7), a data plane DMZ tier **948** (e.g. the data plane DMZ tier **748** of FIG. 7), and a data plane data tier **950** (e.g. the data plane data tier **750** of FIG. 7). The data plane DMZ tier **948** can include LB subnet(s) **922** that can be communicatively coupled to trusted app subnet(s) **960** and untrusted app subnet(s) **962** of the data plane app tier **946** and the Internet gateway **934** contained in the data plane VCN **918**. The trusted app subnet(s) **960** can be communicatively coupled to the service gateway **936** contained in the data plane VCN **918**, the NAT gateway **938** contained in the data plane VCN **918**, and DB subnet(s) **930** contained in the data plane data tier **950**. The untrusted app subnet(s) **962** can be communicatively coupled to the service gateway **936** contained in the data plane VCN **918** and DB subnet(s) **930** contained in the data plane data tier **950**. The data plane data tier **950** can include DB subnet(s) **930** that can be communicatively coupled to the service gateway **936** contained in the data plane VCN **918**.

(94) The untrusted app subnet(s) **962** can include one or more primary VNICs **964(1)-(N)** that can be communicatively coupled to tenant virtual machines (VMs) **966(1)-(N)**. Each tenant VM **966(1)-(N)** can be communicatively coupled to a respective app subnet **967(1)-(N)** that can be contained in respective container egress VCNs **968(1)-(N)** that can be contained in respective customer tenancies **970(1)-(N)**. Respective secondary VNICs **972(1)-(N)** can facilitate communication between the untrusted app subnet(s) **962** contained in the data plane VCN **918** and the app subnet contained in the container egress VCNs **968(1)-(N)**. Each container egress VCNs **968(1)-(N)** can include a NAT gateway **938** that can be communicatively coupled to public Internet **954** (e.g. public Internet **754** of FIG. 7).

(95) The Internet gateway **934** contained in the control plane VCN **916** and contained in the data plane VCN **918** can be communicatively coupled to a metadata management service **952** (e.g. the metadata management system **752** of FIG. 7) that can be communicatively coupled to public Internet **954**. Public Internet **954** can be communicatively coupled to the NAT gateway **938** contained in the control plane VCN **916** and contained in the data plane VCN **918**. The service gateway **936** contained in the control plane VCN **916** and contained in the data plane VCN **918** can be communicatively couple to cloud services **956**.

(96) In some embodiments, the data plane VCN **918** can be integrated with customer tenancies **970**. This integration can be useful or desirable for customers of the IaaS provider in some cases such as a case that may desire support when executing code. The customer may provide code to run that may be destructive, may communicate with other customer resources, or may otherwise cause undesirable effects. In response to this, the IaaS provider may determine whether to run code given to the IaaS provider by the customer.

(97) In some examples, the customer of the IaaS provider may grant temporary network access to the IaaS provider and request a function to be attached to the data plane tier app **946**. Code to run the function may be executed in the VMs **966(1)-(N)**, and the code may not be configured to run anywhere else on the data plane VCN **918**. Each VM **966(1)-(N)** may be connected to one customer tenancy **970**. Respective containers **971(1)-(N)** contained in the VMs **966(1)-(N)** may be

configured to run the code. In this case, there can be a dual isolation (e.g., the containers **971(1)-(N)** running code, where the containers **971(1)-(N)** may be contained in at least the VM **966(1)-(N)** that are contained in the untrusted app subnet(s) **962**), which may help prevent incorrect or otherwise undesirable code from damaging the network of the IaaS provider or from damaging a network of a different customer. The containers **971(1)-(N)** may be communicatively coupled to the customer tenancy **970** and may be configured to transmit or receive data from the customer tenancy **970**. The containers **971(1)-(N)** may not be configured to transmit or receive data from any other entity in the data plane VCN **918**. Upon completion of running the code, the IaaS provider may kill or otherwise dispose of the containers **971(1)-(N)**.

(98) In some embodiments, the trusted app subnet(s) **960** may run code that may be owned or operated by the IaaS provider. In this embodiment, the trusted app subnet(s) **960** may be communicatively coupled to the DB subnet(s) **930** and be configured to execute CRUD operations in the DB subnet(s) **930**. The untrusted app subnet(s) **962** may be communicatively coupled to the DB subnet(s) **930**, but in this embodiment, the untrusted app subnet(s) may be configured to execute read operations in the DB subnet(s) **930**. The containers **971(1)-(N)** that can be contained in the VM **966(1)-(N)** of each customer and that may run code from the customer may not be communicatively coupled with the DB subnet(s) **930**.

(99) In other embodiments, the control plane VCN **916** and the data plane VCN **918** may not be directly communicatively coupled. In this embodiment, there may be no direct communication between the control plane VCN **916** and the data plane VCN **918**. However, communication can occur indirectly through at least one method. An LPG **910** may be established by the IaaS provider that can facilitate communication between the control plane VCN **916** and the data plane VCN **918**. In another example, the control plane VCN **916** or the data plane VCN **918** can make a call to cloud services **956** via the service gateway **936**. For example, a call to cloud services **956** from the control plane VCN **916** can include a request for a service that can communicate with the data plane VCN **918**.

(100) FIG. **10** is a block diagram **1000** illustrating another example pattern of an IaaS architecture, according to at least one embodiment. Service operators **1002** (e.g. service operators **702** of FIG. **7**) can be communicatively coupled to a secure host tenancy **1004** (e.g. the secure host tenancy **704** of FIG. **7**) that can include a virtual cloud network (VCN) **1006** (e.g. the VCN **706** of FIG. **7**) and a secure host subnet **1008** (e.g. the secure host subnet **708** of FIG. **7**). The VCN **1006** can include an LPG **1010** (e.g. the LPG **710** of FIG. **7**) that can be communicatively coupled to an SSH VCN **1012** (e.g. the SSH VCN **712** of FIG. **7**) via an LPG **1010** contained in the SSH VCN **1012**. The SSH VCN **1012** can include an SSH subnet **1014** (e.g. the SSH subnet **714** of FIG. **7**), and the SSH VCN **1012** can be communicatively coupled to a control plane VCN **1016** (e.g. the control plane VCN **716** of FIG. **7**) via an LPG **1010** contained in the control plane VCN **1016** and to a data plane VCN **1018** (e.g. the data plane **718** of FIG. **7**) via an LPG **1010** contained in the data plane VCN **1018**. The control plane VCN **1016** and the data plane VCN **1018** can be contained in a service tenancy **1019** (e.g. the service tenancy **719** of FIG. **7**).

(101) The control plane VCN **1016** can include a control plane DMZ tier **1020** (e.g. the control plane DMZ tier **720** of FIG. **7**) that can include LB subnet(s) **1022** (e.g. LB subnet(s) **722** of FIG. **7**), a control plane app tier **1024** (e.g. the control plane app tier **724** of FIG. **7**) that can include app subnet(s) **1026** (e.g. app subnet(s) **726** of FIG. **7**), a control plane data tier **1028** (e.g. the control plane data tier **728** of FIG. **7**) that can include DB subnet(s) **1030** (e.g. DB subnet(s) **930** of FIG. **9**). The LB subnet(s) **1022** contained in the control plane DMZ tier **1020** can be communicatively coupled to the app subnet(s) **1026** contained in the control plane app tier **1024** and to an Internet gateway **1034** (e.g. the Internet gateway **734** of FIG. **7**) that can be contained in the control plane VCN **1016**, and the app subnet(s) **1026** can be communicatively coupled to the DB subnet(s) **1030** contained in the control plane data tier **1028** and to a service gateway **1036** (e.g. the service gateway of FIG. **7**) and a network address translation (NAT) gateway **1038** (e.g. the NAT gateway

738 of FIG. 7). The control plane VCN **1016** can include the service gateway **1036** and the NAT gateway **1038**.

(102) The data plane VCN **1018** can include a data plane app tier **1046** (e.g. the data plane app tier **746** of FIG. 7), a data plane DMZ tier **1048** (e.g. the data plane DMZ tier **748** of FIG. 7), and a data plane data tier **1050** (e.g. the data plane data tier **750** of FIG. 7). The data plane DMZ tier **1048** can include LB subnet(s) **1022** that can be communicatively coupled to trusted app subnet(s) **1060** (e.g. trusted app subnet(s) **960** of FIG. 9) and untrusted app subnet(s) **1062** (e.g. untrusted app subnet(s) **962** of FIG. 9) of the data plane app tier **1046** and the Internet gateway **1034** contained in the data plane VCN **1018**. The trusted app subnet(s) **1060** can be communicatively coupled to the service gateway **1036** contained in the data plane VCN **1018**, the NAT gateway **1038** contained in the data plane VCN **1018**, and DB subnet(s) **1030** contained in the data plane data tier **1050**. The untrusted app subnet(s) **1062** can be communicatively coupled to the service gateway **1036** contained in the data plane VCN **1018** and DB subnet(s) **1030** contained in the data plane data tier **1050**. The data plane data tier **1050** can include DB subnet(s) **1030** that can be communicatively coupled to the service gateway **1036** contained in the data plane VCN **1018**.

(103) The untrusted app subnet(s) **1062** can include primary VNICs **1064(1)-(N)** that can be communicatively coupled to tenant virtual machines (VMs) **1066(1)-(N)** residing within the untrusted app subnet(s) **1062**. Each tenant VM **1066(1)-(N)** can run code in a respective container **1067(1)-(N)**, and be communicatively coupled to an app subnet **1026** that can be contained in a data plane app tier **1046** that can be contained in a container egress VCN **1068**. Respective secondary VNICs **1072(1)-(N)** can facilitate communication between the untrusted app subnet(s) **1062** contained in the data plane VCN **1018** and the app subnet contained in the container egress VCN **1068**. The container egress VCN can include a NAT gateway **1038** that can be communicatively coupled to public Internet **1054** (e.g. public Internet **754** of FIG. 7).

(104) The Internet gateway **1034** contained in the control plane VCN **1016** and contained in the data plane VCN **1018** can be communicatively coupled to a metadata management service **1052** (e.g. the metadata management system **752** of FIG. 7) that can be communicatively coupled to public Internet **1054**. Public Internet **1054** can be communicatively coupled to the NAT gateway **1038** contained in the control plane VCN **1016** and contained in the data plane VCN **1018**. The service gateway **1036** contained in the control plane VCN **1016** and contained in the data plane VCN **1018** can be communicatively couple to cloud services **1056**.

(105) In some examples, the pattern illustrated by the architecture of block diagram **1000** of FIG. **10** may be considered an exception to the pattern illustrated by the architecture of block diagram **900** of FIG. **9** and may be desirable for a customer of the IaaS provider if the IaaS provider cannot directly communicate with the customer (e.g., a disconnected region). The respective containers **1067(1)-(N)** that are contained in the VMs **1066(1)-(N)** for each customer can be accessed in real-time by the customer. The containers **1067(1)-(N)** may be configured to make calls to respective secondary VNICs **1072(1)-(N)** contained in app subnet(s) **1026** of the data plane app tier **1046** that can be contained in the container egress VCN **1068**. The secondary VNICs **1072(1)-(N)** can transmit the calls to the NAT gateway **1038** that may transmit the calls to public Internet **1054**. In this example, the containers **1067(1)-(N)** that can be accessed in real-time by the customer can be isolated from the control plane VCN **1016** and can be isolated from other entities contained in the data plane VCN **1018**. The containers **1067(1)-(N)** may also be isolated from resources from other customers.

(106) In other examples, the customer can use the containers **1067(1)-(N)** to call cloud services **1056**. In this example, the customer may run code in the containers **1067(1)-(N)** that requests a service from cloud services **1056**. The containers **1067(1)-(N)** can transmit this request to the secondary VNICs **1072(1)-(N)** that can transmit the request to the NAT gateway that can transmit the request to public Internet **1054**. Public Internet **1054** can transmit the request to LB subnet(s) **1022** contained in the control plane VCN **1016** via the Internet gateway **1034**. In response to

determining the request is valid, the LB subnet(s) can transmit the request to app subnet(s) **1026** that can transmit the request to cloud services **1056** via the service gateway **1036**.

(107) It should be appreciated that IaaS architectures **700, 800, 900, 1000** depicted in the figures may have other components than those depicted. Further, the embodiments shown in the figures are only some examples of a cloud infrastructure system that may incorporate an embodiment of the disclosure. In some other embodiments, the IaaS systems may have more or fewer components than shown in the figures, may combine two or more components, or may have a different configuration or arrangement of components.

(108) In certain embodiments, the IaaS systems described herein may include a suite of applications, middleware, and database service offerings that are delivered to a customer in a self-service, subscription-based, elastically scalable, reliable, highly available, and secure manner. An example of such an IaaS system is the Oracle Cloud Infrastructure (OCI) provided by the present assignee.

(109) FIG. **11** illustrates an example computer system **1100**, in which various embodiments may be implemented. The system **1100** may be used to implement any of the computer systems described above. As shown in the figure, computer system **1100** includes a processing unit **1104** that communicates with a number of peripheral subsystems via a bus subsystem **1102**. These peripheral subsystems may include a processing acceleration unit **1106**, an I/O subsystem **1108**, a storage subsystem **1118** and a communications subsystem **1124**. Storage subsystem **1118** includes tangible computer-readable storage media **1122** and a system memory **1110**.

(110) Bus subsystem **1102** provides a mechanism for letting the various components and subsystems of computer system **1100** communicate with each other as intended. Although bus subsystem **1102** is shown schematically as a single bus, alternative embodiments of the bus subsystem may utilize multiple buses. Bus subsystem **1102** may be any of several types of bus structures including a memory bus or memory controller, a peripheral bus, and a local bus using any of a variety of bus architectures. For example, such architectures may include an Industry Standard Architecture (ISA) bus, Micro Channel Architecture (MCA) bus, Enhanced ISA (EISA) bus, Video Electronics Standards Association (VESA) local bus, and Peripheral Component Interconnect (PCI) bus, which can be implemented as a Mezzanine bus manufactured to the IEEE P1386.1 standard.

(111) Processing unit **1104**, which can be implemented as one or more integrated circuits (e.g., a conventional microprocessor or microcontroller), controls the operation of computer system **1100**. One or more processors may be included in processing unit **1104**. These processors may include single core or multicore processors. In certain embodiments, processing unit **1104** may be implemented as one or more independent processing units **1132** and/or **1134** with single or multicore processors included in each processing unit. In other embodiments, processing unit **1104** may also be implemented as a quad-core processing unit formed by integrating two dual-core processors into a single chip.

(112) In various embodiments, processing unit **1104** can execute a variety of programs in response to program code and can maintain multiple concurrently executing programs or processes. At any given time, some or all of the program code to be executed can be resident in processor(s) **1104** and/or in storage subsystem **1118**. Through suitable programming, processor(s) **1104** can provide various functionalities described above. Computer system **1100** may additionally include a processing acceleration unit **1106**, which can include a digital signal processor (DSP), a special-purpose processor, and/or the like.

(113) I/O subsystem **1108** may include user interface input devices and user interface output devices. User interface input devices may include a keyboard, pointing devices such as a mouse or trackball, a touchpad or touch screen incorporated into a display, a scroll wheel, a click wheel, a dial, a button, a switch, a keypad, audio input devices with voice command recognition systems, microphones, and other types of input devices. User interface input devices may include, for

example, motion sensing and/or gesture recognition devices such as the Microsoft Kinect® motion sensor that enables users to control and interact with an input device, such as the Microsoft Xbox® 360 game controller, through a natural user interface using gestures and spoken commands. User interface input devices may also include eye gesture recognition devices such as the Google Glass® blink detector that detects eye activity (e.g., ‘blinking’ while taking pictures and/or making a menu selection) from users and transforms the eye gestures as input into an input device (e.g., Google Glass®). Additionally, user interface input devices may include voice recognition sensing devices that enable users to interact with voice recognition systems (e.g., Siri® navigator), through voice commands.

(114) User interface input devices may also include, without limitation, three dimensional (3D) mice, joysticks or pointing sticks, gamepads and graphic tablets, and audio/visual devices such as speakers, digital cameras, digital camcorders, portable media players, webcams, image scanners, fingerprint scanners, barcode reader 3D scanners, 3D printers, laser rangefinders, and eye gaze tracking devices. Additionally, user interface input devices may include, for example, medical imaging input devices such as computed tomography, magnetic resonance imaging, position emission tomography, medical ultrasonography devices. User interface input devices may also include, for example, audio input devices such as MIDI keyboards, digital musical instruments and the like.

(115) User interface output devices may include a display subsystem, indicator lights, or non-visual displays such as audio output devices, etc. The display subsystem may be a cathode ray tube (CRT), a flat-panel device, such as that using a liquid crystal display (LCD) or plasma display, a projection device, a touch screen, and the like. In general, use of the term “output device” is intended to include all possible types of devices and mechanisms for outputting information from computer system **1100** to a user or other computer. For example, user interface output devices may include, without limitation, a variety of display devices that visually convey text, graphics and audio/video information such as monitors, printers, speakers, headphones, automotive navigation systems, plotters, voice output devices, and modems.

(116) Computer system **1100** may comprise a storage subsystem **1118** that comprises software elements, shown as being currently located within a system memory **1110**. System memory **1110** may store program instructions that are loadable and executable on processing unit **1104**, as well as data generated during the execution of these programs.

(117) Depending on the configuration and type of computer system **1100**, system memory **1110** may be volatile (such as random access memory (RAM)) and/or non-volatile (such as read-only memory (ROM), flash memory, etc.) The RAM typically contains data and/or program modules that are immediately accessible to and/or presently being operated and executed by processing unit **1104**. In some implementations, system memory **1110** may include multiple different types of memory, such as static random access memory (SRAM) or dynamic random access memory (DRAM). In some implementations, a basic input/output system (BIOS), containing the basic routines that help to transfer information between elements within computer system **1100**, such as during start-up, may typically be stored in the ROM. By way of example, and not limitation, system memory **1110** also illustrates application programs **1112**, which may include client applications, Web browsers, mid-tier applications, relational database management systems (RDBMS), etc., program data **1114**, and an operating system **1116**. By way of example, operating system **1116** may include various versions of Microsoft Windows®, Apple Macintosh®, and/or Linux operating systems, a variety of commercially-available UNIX® or UNIX-like operating systems (including without limitation the variety of GNU/Linux operating systems, the Google Chrome® OS, and the like) and/or mobile operating systems such as iOS, Windows® Phone, Android® OS, BlackBerry® 11 OS, and Palm® OS operating systems.

(118) Storage subsystem **1118** may also provide a tangible computer-readable storage medium for storing the basic programming and data constructs that provide the functionality of some

embodiments. Software (programs, code modules, instructions) that when executed by a processor provide the functionality described above may be stored in storage subsystem **1118**. These software modules or instructions may be executed by processing unit **1104**. Storage subsystem **1118** may also provide a repository for storing data used in accordance with the present disclosure.

(119) Storage subsystem **1100** may also include a computer-readable storage media reader **1120** that can further be connected to computer-readable storage media **1122**. Together and, optionally, in combination with system memory **1110**, computer-readable storage media **1122** may comprehensively represent remote, local, fixed, and/or removable storage devices plus storage media for temporarily and/or more permanently containing, storing, transmitting, and retrieving computer-readable information.

(120) Computer-readable storage media **1122** containing code, or portions of code, can also include any appropriate media known or used in the art, including storage media and communication media, such as but not limited to, volatile and non-volatile, removable and non-removable media implemented in any method or technology for storage and/or transmission of information. This can include tangible computer-readable storage media such as RAM, ROM, electronically erasable programmable ROM (EEPROM), flash memory or other memory technology, CD-ROM, digital versatile disk (DVD), or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or other tangible computer readable media. This can also include nontangible computer-readable media, such as data signals, data transmissions, or any other medium which can be used to transmit the desired information and which can be accessed by computing system **1100**.

(121) By way of example, computer-readable storage media **1122** may include a hard disk drive that reads from or writes to non-removable, nonvolatile magnetic media, a magnetic disk drive that reads from or writes to a removable, nonvolatile magnetic disk, and an optical disk drive that reads from or writes to a removable, nonvolatile optical disk such as a CD ROM, DVD, and Blu-Ray® disk, or other optical media. Computer-readable storage media **1122** may include, but is not limited to, Zip® drives, flash memory cards, universal serial bus (USB) flash drives, secure digital (SD) cards, DVD disks, digital video tape, and the like. Computer-readable storage media **1122** may also include, solid-state drives (SSD) based on non-volatile memory such as flash-memory based SSDs, enterprise flash drives, solid state ROM, and the like, SSDs based on volatile memory such as solid state RAM, dynamic RAM, static RAM, DRAM-based SSDs, magnetoresistive RAM (MRAM) SSDs, and hybrid SSDs that use a combination of DRAM and flash memory based SSDs. The disk drives and their associated computer-readable media may provide non-volatile storage of computer-readable instructions, data structures, program modules, and other data for computer system **1100**.

(122) Communications subsystem **1124** provides an interface to other computer systems and networks. Communications subsystem **1124** serves as an interface for receiving data from and transmitting data to other systems from computer system **1100**. For example, communications subsystem **1124** may enable computer system **1100** to connect to one or more devices via the Internet. In some embodiments communications subsystem **1124** can include radio frequency (RF) transceiver components for accessing wireless voice and/or data networks (e.g., using cellular telephone technology, advanced data network technology, such as 3G, 4G or EDGE (enhanced data rates for global evolution), WiFi (IEEE 802.11 family standards, or other mobile communication technologies, or any combination thereof), global positioning system (GPS) receiver components, and/or other components. In some embodiments communications subsystem **1124** can provide wired network connectivity (e.g., Ethernet) in addition to or instead of a wireless interface.

(123) In some embodiments, communications subsystem **1124** may also receive input communication in the form of structured and/or unstructured data feeds **1126**, event streams **1128**, event updates **1130**, and the like on behalf of one or more users who may use computer system **1100**.

(124) By way of example, communications subsystem **1124** may be configured to receive data

feeds **1126** in real-time from users of social networks and/or other communication services such as Twitter® feeds, Facebook® updates, web feeds such as Rich Site Summary (RSS) feeds, and/or real-time updates from one or more third party information sources.

(125) Additionally, communications subsystem **1124** may also be configured to receive data in the form of continuous data streams, which may include event streams **1128** of real-time events and/or event updates **1130**, that may be continuous or unbounded in nature with no explicit end. Examples of applications that generate continuous data may include, for example, sensor data applications, financial tickers, network performance measuring tools (e.g. network monitoring and traffic management applications), clickstream analysis tools, automobile traffic monitoring, and the like.

(126) Communications subsystem **1124** may also be configured to output the structured and/or unstructured data feeds **1126**, event streams **1128**, event updates **1130**, and the like to one or more databases that may be in communication with one or more streaming data source computers coupled to computer system **1100**.

(127) Computer system **1100** can be one of various types, including a handheld portable device (e.g., an iPhone® cellular phone, an iPad® computing tablet, a PDA), a wearable device (e.g., a Google Glass® head mounted display), a PC, a workstation, a mainframe, a kiosk, a server rack, or any other data processing system.

(128) Due to the ever-changing nature of computers and networks, the description of computer system **1100** depicted in the figure is intended only as a specific example. Many other configurations having more or fewer components than the system depicted in the figure are possible. For example, customized hardware might also be used and/or particular elements might be implemented in hardware, firmware, software (including applets), or a combination. Further, connection to other computing devices, such as network input/output devices, may be employed. Based on the disclosure and teachings provided herein, a person of ordinary skill in the art will appreciate other ways and/or methods to implement the various embodiments.

(129) Although specific embodiments have been described, various modifications, alterations, alternative constructions, and equivalents are also encompassed within the scope of the disclosure. Embodiments are not restricted to operation within certain specific data processing environments, but are free to operate within a plurality of data processing environments. Additionally, although embodiments have been described using a particular series of transactions and steps, it should be apparent to those skilled in the art that the scope of the present disclosure is not limited to the described series of transactions and steps. Various features and aspects of the above-described embodiments may be used individually or jointly.

(130) Further, while embodiments have been described using a particular combination of hardware and software, it should be recognized that other combinations of hardware and software are also within the scope of the present disclosure. Embodiments may be implemented only in hardware, or only in software, or using combinations thereof. The various processes described herein can be implemented on the same processor or different processors in any combination. Accordingly, where components or modules are described as being configured to perform certain operations, such configuration can be accomplished, e.g., by designing electronic circuits to perform the operation, by programming programmable electronic circuits (such as microprocessors) to perform the operation, or any combination thereof. Processes can communicate using a variety of techniques including but not limited to conventional techniques for inter process communication, and different pairs of processes may use different techniques, or the same pair of processes may use different techniques at different times.

(131) The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense. It will, however, be evident that additions, subtractions, deletions, and other modifications and changes may be made thereunto without departing from the broader spirit and scope as set forth in the claims. Thus, although specific disclosure embodiments have been described, these are not intended to be limiting. Various modifications and equivalents are within

the scope of the following claims.

(132) The use of the terms “a” and “an” and “the” and similar referents in the context of describing the disclosed embodiments (especially in the context of the following claims) are to be construed to cover both the singular and the plural, unless otherwise indicated herein or clearly contradicted by context. The terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (i.e., meaning “including, but not limited to,”) unless otherwise noted. The term “connected” is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within the range, unless otherwise indicated herein and each separate value is incorporated into the specification as if it were individually recited herein. All methods described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. The use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments and does not pose a limitation on the scope of the disclosure unless otherwise claimed. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the disclosure.

(133) Disjunctive language such as the phrase “at least one of X, Y, or Z,” unless specifically stated otherwise, is intended to be understood within the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present.

(134) Preferred embodiments of this disclosure are described herein, including the best mode known for carrying out the disclosure. Variations of those preferred embodiments may become apparent to those of ordinary skill in the art upon reading the foregoing description. Those of ordinary skill should be able to employ such variations as appropriate and the disclosure may be practiced otherwise than as specifically described herein. Accordingly, this disclosure includes all modifications and equivalents of the subject matter recited in the claims appended hereto as permitted by applicable law. Moreover, any combination of the above-described elements in all possible variations thereof is encompassed by the disclosure unless otherwise indicated herein.

(135) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to the same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(136) In the foregoing specification, aspects of the disclosure are described with reference to specific embodiments thereof, but those skilled in the art will recognize that the disclosure is not limited thereto. Various features and aspects of the above-described disclosure may be used individually or jointly. Further, embodiments can be utilized in any number of environments and applications beyond those described herein without departing from the broader spirit and scope of the specification. The specification and drawings are, accordingly, to be regarded as illustrative rather than restrictive.

Claims

1. One or more non-transitory computer-readable media storing instructions that, when executed by one or more hardware processors, cause performance of operations comprising: selecting a first zone segment for a defragmentation process, the first zone segment comprising a first plurality of zones, wherein the first zone segment comprises stored data that has been written to the first zone segment via a prior write operation, wherein a configuration associated with the first zone segment results in the stored data impeding any writing of additional data to unused storage capacity of a first subset of the first plurality of zones of the first zone segment via any new write operation while the stored data is maintained within one or more source zones of a second subset of the first

plurality of zones of the first zone segment, wherein the first subset of the first plurality of zones precedes the second subset of the first plurality of zones within the first zone segment; accessing the stored data comprised within the one or more source zones of the first zone segment; writing, via a single request specifying a write operation command, the stored data to a plurality of target zones comprised within a second zone segment, wherein the plurality of target zones comprises at least (a) a first zone and (b) a second zone that is subsequent to the first zone in the plurality of target zones; and resetting at least one of the one or more source zones of the first zone segment to an empty status or an open status.

2. The one or more non-transitory computer-readable media of claim 1, wherein the plurality of target zones is comprised within a second plurality of sequentially-organized zones of a second zone segment, and wherein writing the stored data to the plurality of target zones of the second zone segment comprises: accessing a segment pointer of the second zone segment that indicates a write location in the second zone segment; writing, starting at the write location, the stored data to the plurality of target zones of the second zone segment; and updating the segment pointer of the second zone segment according to a data endpoint of the stored data in the second zone segment.

3. The one or more non-transitory computer-readable media of claim 1, wherein the first zone segment comprises a first plurality of sequentially-organized zones, and wherein accessing the stored data comprised within the one or more source zones of the first zone segment comprises: reading, in a single read operation, from at least a third zone and a fourth zone based on the fourth zone being organized subsequent to the third zone in first zone segment.

4. The one or more non-transitory computer-readable media of claim 1, wherein selecting the first zone segment for the defragmentation process comprises: determining a usage fraction of the first zone segment based on volume metadata associated with the first zone segment, wherein the usage fraction of the first zone segment indicates a portion of a storage capacity of the first zone segment that is mapped to the volume metadata associated with the first zone segment; comparing the usage fraction of the first zone segment to a usage threshold; and determining that the usage fraction of the first zone segment does not satisfy the usage threshold.

5. The one or more non-transitory computer-readable media of claim 1, wherein the stored data comprised within the one or more source zones of the first zone segment is accessed based on (a) volume metadata associated with the stored data and/or (b) scanning the first zone segment.

6. The one or more non-transitory computer-readable media of claim 1, wherein the operations further comprise: subsequent to accessing the stored data comprised within the one or more source zones of the first zone segment: remapping the stored data from the one or more source zones to the plurality of target zones.

7. The one or more non-transitory computer-readable media of claim 6, wherein remapping the stored data from the one or more source zones to the plurality of target zones comprises: modifying volume metadata associated with the stored data such that the volume metadata associated with the stored data is mapped to the plurality of target zones.

8. The one or more non-transitory computer-readable media of claim 1, wherein resetting at least one of the one or more source zones of the first zone segment comprises: (a) changing a status of a source zone of the one or more source zones to an open status or empty status and/or (b) altering volume metadata to remove mapping information associated with the source zone of the one or more sources zones.

9. The one or more non-transitory computer-readable media of claim 1, wherein the operations further comprise: prior to writing the stored data to the plurality of target zones: writing the stored data to buffer storage, the buffer storage comprising a second zone segment and/or random access memory (RAM); and mapping the stored data to the buffer storage.

10. A method, comprising: selecting a first zone segment for a defragmentation process, the first zone segment comprising a first plurality of zones, wherein the first zone segment comprises stored data that has been written to the first zone segment via a prior write operation, wherein a

configuration associated with the first zone segment results in the stored data impeding any writing of additional data to unused storage capacity of a first subset of the first plurality of zones of the first zone segment via any new write operation while the stored data is maintained within one or more source zones of a second subset of the first plurality of zones of the first zone segment, wherein the first subset of the first plurality of zones precedes the second subset of the first plurality of zones within the first zone segment; accessing the stored data comprised within the one or more source zones of the first zone segment; writing, via a single request specifying a write operation command, the stored data to a plurality of target zones comprised within a second zone segment, wherein the plurality of target zones comprises at least (a) a first zone and (b) a second zone that is subsequent to the first zone in the plurality of target zones; and resetting at least one of the one or more source zones of the first zone segment to an empty status or an open status wherein the method is performed by at least one device including a hardware processor.

11. The method of claim 10, wherein the plurality of target zones is comprised within a second plurality of sequentially-organized zones of a second zone segment, and wherein writing the stored data to the plurality of target zones of the second zone segment comprises: accessing a segment pointer of the second zone segment that indicates a write location in the second zone segment; writing, starting at the write location, the stored data to the plurality of target zones of the second zone segment; and updating the segment pointer of the second zone segment according to a data endpoint of the stored data in the second zone segment.

12. The method of claim 10, wherein the first zone segment comprises a first plurality of sequentially-organized zones, and wherein accessing the stored data comprised within the one or more source zones of the first zone segment comprises: reading, in a single read operation, from at least a third zone and a fourth zone based on the fourth zone being organized subsequent to the third zone in first zone segment.

13. The method of claim 10, wherein selecting the first zone segment for the defragmentation process comprises: determining a usage fraction of the first zone segment based on volume metadata associated with the first zone segment, wherein the usage fraction of the first zone segment indicates a portion of a storage capacity of the first zone segment that is mapped to the volume metadata associated with the first zone segment; comparing the usage fraction of the first zone segment to a usage threshold; and determining that the usage fraction of the first zone segment does not satisfy the usage threshold.

14. The method of claim 10, wherein the stored data comprised within the one or more source zones of the first zone segment is accessed based on (a) volume metadata associated with the stored data and/or (b) scanning the first zone segment.

15. The method of claim 10, further comprising: subsequent to accessing the stored data comprised within the one or more source zones of the first zone segment: remapping the stored data from the one or more source zones to the plurality of target zones.

16. The method of claim 15, wherein remapping the stored data from the one or more source zones to the plurality of target zones comprises: modifying volume metadata associated with the stored data such that the volume metadata associated with the stored data is mapped to the plurality of target zones.

17. The method of claim 10, wherein the operations further comprise: prior to writing the stored data to the plurality of target zones: writing the stored data to buffer storage, the buffer storage comprising a second zone segment and/or random access memory (RAM); and mapping the stored data to the buffer storage.

18. A system, comprising: at least one device including a hardware processor; and the system being configured to perform operations comprising: selecting a first zone segment for a defragmentation process, the first zone segment comprising a first plurality of zones, wherein the first zone segment comprises stored data that has been written to the first zone segment via a prior write operation, wherein a configuration associated with the first zone segment results in the stored data impeding

any writing of additional data to unused storage capacity of a first subset of the first plurality of zones of the first zone segment via any new write operation while the stored data is maintained within one or more source zones of a second subset of the first plurality of zones of the first zone segment, wherein the first subset of the first plurality of zones precedes the second subset of the first plurality of zones within the first zone segment; accessing the stored data comprised within the one or more source zones of the first zone segment; writing, via a single request specifying a write operation command, the stored data to a plurality of target zones comprised within a second zone segment, wherein the plurality of target zones comprises at least (a) a first zone and (b) a second zone that is subsequent to the first zone in the plurality of target zones; and resetting at least one of the one or more source zones of the first zone segment to an empty status or an open status.

19. The system of claim 18, wherein the plurality of target zones is comprised within a second plurality of sequentially-organized zones of a second zone segment, and wherein writing the stored data to the plurality of target zones of the second zone segment comprises: accessing a segment pointer of the second zone segment that indicates a write location in the second zone segment; writing, starting at the write location, the stored data to the plurality of target zones of the second zone segment; and updating the segment pointer of the second zone segment according to a data endpoint of the stored data in the second zone segment.

20. The system of claim 18, wherein the first zone segment comprises a first plurality of sequentially-organized zones, and wherein accessing the stored data comprised within the one or more source zones of the first zone segment comprises: reading, in a single read operation, from at least a third zone and a fourth zone based on the fourth zone being organized subsequent to the third zone in first zone segment.
