

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386543
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Cariello; Giuseppe et al.

Opportunistic storage of non-write-boosted data in write booster cache memory

Abstract

In some implementations, a memory device may receive a write command that includes data to be written to the memory device. The memory device may receive an indication that single-level cell data caching is deactivated for the data. The memory device may determine whether the data is associated with a first data type or a second data type. The memory device may selectively write the data to single-level cell cache memory or multi-level cell main memory based on a determination of whether the data is associated with the first data type or the second data type and a determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching.

Inventors: Cariello; Giuseppe (Boise, ID), Parry; Jonathan S. (Boise, ID), Basu; Reshmi (Boise, ID)

Applicant: Micron Technology, Inc. (Boise, ID)

Family ID: 91473763

Assignee: Micron Technology, Inc. (Boise, ID)

Appl. No.: 18/505661

Filed: November 09, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240201888 A1	Jun. 20, 2024

Related U.S. Application Data

us-provisional-application US 63387536 20221215

Publication Classification

Int. Cl.: G06F12/00 (20060101); G06F3/06 (20060101)

U.S. Cl.:

CPC G06F3/065 (20130101); G06F3/0608 (20130101); G06F3/064 (20130101); G06F3/0679 (20130101);

Field of Classification Search

CPC: G06F (3/0608); G06F (3/064); G06F (3/0613); G06F (3/0659); G06F (3/065); G06F (3/0679); G06F (12/02); G06F (12/084); G06F (12/0246); G06F (2212/1032); G06F (2212/2022); G06F (2212/601); G06F (2212/7204); G06F (2212/7205); G06F (2212/7206); G06F (2212/7211)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2014/0337566	12/2013	Oh	711/103	G06F 3/0659
2021/0365202	12/2020	Shin	N/A	G06F 3/0679

Primary Examiner: Gu; Shawn X

Attorney, Agent or Firm: Harrity & Harrity, LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This Patent application claims priority to U.S. Provisional Patent Application No. 63/387,536, filed on Dec. 15, 2022, and entitled “OPPORTUNISTIC STORAGE OF NON-WRITE-BOOSTED DATA IN WRITE BOOSTER CACHE MEMORY.” The disclosure of the prior Application is considered part of and is incorporated by reference into this Patent Application.

TECHNICAL FIELD

(1) The present disclosure generally relates to memory devices, memory device operations, and, for example, to opportunistic storage of non-write-boosted data in write booster cache memory.

BACKGROUND

(2) Memory devices are widely used to store information in various electronic devices. A memory device includes memory cells. A memory cell is an electronic circuit capable of being programmed to a data state of two or more data states. For example, a memory cell may be programmed to a data state that represents a single binary value, often denoted by a binary “1” or a binary “0.” As another example, a memory cell may be programmed to a data state that represents a fractional value (e.g., 0.5, 1.5, or the like). To store information, an electronic device may write to, or program, a set of memory cells. To access the stored information, the electronic device may read, or sense, the stored state from the set of memory cells.

(3) Various types of memory devices exist, including random access memory (RAM), read only

memory (ROM), dynamic RAM (DRAM), static RAM (SRAM), synchronous dynamic RAM (SDRAM), ferroelectric RAM (FeRAM), magnetic RAM (MRAM), resistive RAM (RRAM), holographic RAM (HRAM), flash memory (e.g., NAND memory and NOR memory), and others. A memory device may be volatile or non-volatile. Non-volatile memory (e.g., flash memory) can store data for extended periods of time even in the absence of an external power source. Volatile memory (e.g., DRAM) may lose stored data over time unless the volatile memory is refreshed by a power source.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a diagram illustrating an example system capable of opportunistic storage of non-write-boosted data in write booster cache memory.
- (2) FIG. 2 is a diagram of example components included in a memory device.
- (3) FIG. 3 is a diagram illustrating an example memory architecture that may be used by the memory device.
- (4) FIG. 4 is a diagram of an example of reserving a portion of write booster cache memory for write-boosted data.
- (5) FIG. 5 is a diagram of an example of opportunistic storage of non-write-boosted data in write booster cache memory.
- (6) FIG. 6 is a diagram of an example of performing a garbage collection operation on write booster cache memory.
- (7) FIGS. 7-9 are flowcharts of example methods associated with opportunistic storage of non-write-boosted data in write booster cache memory.

DETAILED DESCRIPTION

(8) Some memory devices may use a technique called “write boosting” (sometimes called “single-level cell data caching” or “write boost data caching” herein) to increase a speed at which data is written to memory. Using write boosting, a memory device may write data to a single-level cell (SLC) non-volatile “write booster” cache memory rather than writing the data to a multi-level cell (e.g., dual-level cell, triple-level cell, quad-level cell or beyond) non-volatile main memory. Because data can be written to SLC cache memory faster than the data can be written to multi-level cell (MLC) main memory, using write boosting increases write speeds. A specific write boosting technique is defined by the Joint Electron Device Engineering Council (JEDEC) universal flash storage (UFS) standard JESD220.

(9) To use write boosting, a host device explicitly instructs the memory device to activate write boosting, such as by transmitting a message to the memory device (e.g., in a write command or prior to transmitting the write command) that instructs the memory device to activate write boosting for one or more write commands. For example, the host device may activate write boosting for a large file that is downloaded by the host device and stored in the memory device. If the host device does not instruct the memory device to activate write boosting, then the memory device does not use write boosting.

(10) However, requiring explicit activation for the memory device to write data to write booster cache memory (e.g., SLC cache memory used for write boosting) may result in inefficient utilization of memory resources and poor flexibility for memory block management. For example, if the host device never activates write boosting or activates write boosting infrequently, then write booster cache memory may rarely or never be used. On the other hand, if the memory device is permitted to unconditionally write data to the write booster cache memory, then the write booster cache memory may be unavailable for write boosting requests from the host device and/or the write booster cache memory may wear out earlier.

(11) Some implementations described herein enable the memory device to opportunistically (e.g., conditionally) write non-write-boostered data to write booster cache memory, which increases write speeds and results in better memory resource utilization. In some implementations, the memory device may reserve a portion of the write booster cache memory for storage of data for which write boosting is activated, which increases the likelihood that some write booster cache memory is available when write boosting is requested by the host device.

(12) Furthermore, some implementations described herein limit the type of data permitted to be stored in the write booster cache memory. As an example, the memory device may store hot data, having a relatively short expected lifespan (e.g., as compared to cold data) and/or that is more frequently accessed, erased, and/or overwritten (e.g., as compared to cold data), in the write booster cache memory. This may extend the lifespan of the memory device because otherwise the hot data would be written to MLC main memory and would be marked as invalid soon after (e.g., due to being hot data with a relatively short lifespan), requiring subsequent erasure and wearing out the MLC main memory more quickly due to shorter program/erase (P/E) cycles. In general, MLC memory has a shorter lifespan (e.g., becoming unreliable in fewer P/E cycles) than SLC memory. As a result, writing hot data to SLC memory rather than MLC memory has a less detrimental impact on the lifespan of the memory device compared to writing the hot data to the MLC memory.

(13) Thus, some implementations described herein extend the lifespan of the memory device (e.g., by reducing write amplification of hot data in MLC memory), improve resource utilization of the memory device (e.g., by permitting flexible use of SLC memory that would otherwise be entirely used for write boosting), and/or improve performance of the memory device (e.g., by ensuring that a portion of write booster memory cache is available for write-boosting operations).

(14) FIG. 1 is a diagram illustrating an example system **100** capable of opportunistic storage of non-write-boostered data in write booster cache memory. The system **100** may include one or more devices, apparatuses, and/or components for performing operations described herein. For example, the system **100** may include a host device **110** and a memory device **120**. The memory device **120** may include a controller **130** and memory **140**. The host device **110** may communicate with the memory device **120** (e.g., the controller **130** of the memory device **120**) via a host interface **150**. The controller **130** and the memory **140** may communicate via a memory interface **160**.

(15) The system **100** may be any electronic device configured to store data in memory. For example, the system **100** may be a computer, a mobile phone, a wired or wireless communication device, a network device, a server, a device in a data center, a device in a cloud computing environment, a vehicle (e.g., an automobile or an airplane), and/or an Internet of Things (IOT) device. The host device **110** may include one or more processors configured to execute instructions and store data in the memory **140**. For example, the host device **110** may include a central processing unit (CPU), a graphics processing unit (GPU), a field-programmable gate array (FPGA), an application-specific integrated circuit (ASIC), and/or another type of processing component.

(16) The memory device **120** may be any electronic device or apparatus configured to store data in memory. In some implementations, the memory device **120** may be an electronic device configured to store data persistently in non-volatile memory. For example, the memory device **120** may be a hard drive, a solid-state drive (SSD), a flash memory device (e.g., a NAND flash memory device or a NOR flash memory device), a universal serial bus (USB) thumb drive, a memory card (e.g., a secure digital (SD) card), a secondary storage device, a non-volatile memory express (NVMe) device, and/or an embedded multimedia card (eMMC) device. In this case, the memory **140** may include non-volatile memory configured to maintain stored data after the memory device **120** is powered off. For example, the memory **140** may include NAND memory or NOR memory. In some implementations, the memory **140** may include volatile memory that requires power to maintain stored data and that loses stored data after the memory device **120** is powered off, such as one or more latches and/or random-access memory (RAM), such as dynamic RAM (DRAM) and/or static RAM (SRAM). For example, the volatile memory may cache data read from or to be

written to non-volatile memory, and/or may cache instructions to be executed by the controller **130**.

(17) The controller **130** may be any device configured to communicate with the host device (e.g., via the host interface **150**) and the memory **140** (e.g., via the memory interface **160**). Additionally, or alternatively, the controller **130** may be configured to control operations of the memory device **120** and/or the memory **140**. For example, the controller **130** may include control logic, a memory controller, a system controller, an ASIC, an FPGA, a processor, a microcontroller, and/or one or more processing components. In some implementations, the controller **130** may be a high-level controller, which may communicate directly with the host device **110** and may instruct one or more low-level controllers regarding memory operations to be performed in connection with the memory **140**. In some implementations, the controller **130** may be a low-level controller, which may receive instructions regarding memory operations from a high-level controller that interfaces directly with the host device **110**. As an example, a high-level controller may be an SSD controller, and a low-level controller may be a non-volatile memory controller (e.g., a NAND controller) or a volatile memory controller (e.g., a DRAM controller). In some implementations, a set of operations described herein as being performed by the controller **130** may be performed by a single controller (e.g., the entire set of operations may be performed by a single high-level controller or a single low-level controller). Alternatively, a set of operations described herein as being performed by the controller **130** may be performed by more than one controller (e.g., a first subset of the operations may be performed by a high-level controller and a second subset of the operations may be performed by a low-level controller).

(18) The host interface **150** enables communication between the host device **110** and the memory device **120**. The host interface **150** may include, for example, a Small Computer System Interface (SCSI), a Serial-Attached SCSI (SAS), a Serial Advanced Technology Attachment (SATA) interface, a Peripheral Component Interconnect Express (PCIe) interface, an NVMe interface, a USB interface, a Universal Flash Storage (UFS) interface, and/or an embedded multimedia card (eMMC) interface.

(19) The memory interface **160** enables communication between the memory device **120** and the memory **140**. The memory interface **160** may include a non-volatile memory interface (e.g., for communicating with non-volatile memory), such as a NAND interface or a NOR interface. Additionally, or alternatively, the memory interface **160** may include a volatile memory interface (e.g., for communicating with volatile memory), such as a double data rate (DDR) interface.

(20) In some implementations, the memory device **120** and/or the controller **130** may be configured to receive, from a host device, an indication of a threshold amount of non-volatile cache memory to be reserved for first data for which write boosting is activated; receive, from the host device, a write command that includes second data for which write boosting is deactivated; determine that the second data is hot data; identify a block of the non-volatile cache memory using a write boost cursor based on determining that the second data is hot data, based on the indication of the threshold amount, and despite write boosting being deactivated for the second data; and write the second data to the block of the non-volatile cache memory.

(21) In some implementations, the memory device **120** and/or the controller **130** may be configured to receive a write command that includes data to be written to the memory device; receive an indication that single-level cell data caching is deactivated for the data; determine whether the data is associated with a first data type or a second data type; and selectively write the data to single-level cell cache memory or multi-level cell main memory based on: a determination of whether the data is associated with the first data type or the second data type, and a determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching.

(22) In some implementations, the memory device **120** and/or the controller **130** may be configured to receive a write command that includes data to be written to the memory device; determine that write boost data caching is deactivated for the data; determine that the data is associated with a first

data type; determine that the data is to be written to cache memory, rather than main memory, based on determining that the data is associated with the first data type and despite write boost data caching being deactivated for the data; and write the data to the cache memory based on determining that the data is to be written to the cache memory.

(23) In some implementations, the memory device **120** and/or the controller **130** may be configured to process a received command to write data to a non-volatile memory array, wherein the command is received without an indication that the data is to be written in a single-level cell (SLC) cache memory of the non-volatile memory array; determine whether the data is associated with a first data type; determine whether the SLC cache memory has sufficient memory not reserved for SLC data caching; and write the data to the SLC cache memory based on a determination that the data is associated with the first data type and that the SLC cache memory has sufficient memory not reserved for SLC data caching.

(24) As indicated above, FIG. **1** is provided as an example. Other examples may differ from what is described with regard to FIG. **1**.

(25) FIG. **2** is a diagram of example components included in a memory device **120**. As described above in connection with FIG. **1**, the memory device **120** may include a controller **130** and memory **140**. As shown in FIG. **2**, the memory **140** may include one or more non-volatile memory arrays **205**, such as one or more NAND memory arrays and/or one or more NOR memory arrays. Additionally, or alternatively, the memory **140** may include one or more volatile memory arrays **210**, such as one or more SRAM arrays and/or one or more DRAM arrays. The controller **130** may transmit signals to and receive signals from a non-volatile memory array **205** using a non-volatile memory interface **215**. The controller **130** may transmit signals to and receive signals from a volatile memory array **210** using a volatile memory interface **220**. In some implementations, a single non-volatile memory array **205** is included in a single corresponding die along with a local controller. Multiple dies, each including a non-volatile memory array **205** and a local controller, may be stacked and/or incorporated into an integrated circuit along with an external (or global) controller that communicates with respective local controllers of each die.

(26) The controller **130** may control operations of the memory **140**, such as by executing one or more instructions. For example, the memory device **120** may store one or more instructions in the memory **140** as firmware, and the controller **130** may execute those one or more instructions. Additionally, or alternatively, the controller **130** may receive one or more instructions from the host device **110** via the host interface **150**, and may execute those one or more instructions. In some implementations, a non-transitory computer-readable medium (e.g., volatile memory and/or non-volatile memory) may store a set of instructions (e.g., one or more instructions or code) for execution by the controller **130**. The controller **130** may execute the set of instructions to perform one or more operations or methods described herein. In some implementations, execution of the set of instructions, by the controller **130**, causes the controller **130** and/or the memory device **120** to perform one or more operations or methods described herein. In some implementations, hardwired circuitry is used instead of or in combination with the one or more instructions to perform one or more operations or methods described herein. Additionally, or alternatively, the controller **130** and/or one or more components of the memory device **120** may be configured to perform one or more operations or methods described herein. An instruction is sometimes called a “command.”

(27) For example, the controller **130** may transmit signals to and/or receive signals from the memory **140** based on the one or more instructions, such as to transfer data to (e.g., write or program), to transfer data from (e.g., read), and/or to erase all or a portion of the memory **140** (e.g., one or more memory cells, pages, sub-blocks, blocks, or planes of the memory **140**). Additionally, or alternatively, the controller **130** may be configured to control access to the memory **140** and/or to provide a translation layer between the host device **110** and the memory **140** (e.g., for mapping logical addresses to physical addresses of a memory array). In some implementations, the controller **130** may translate a host interface command (e.g., a command received from the host

device **110**) into a memory interface command (e.g., a command for performing an operation on a memory array).

(28) As shown in FIG. 2, the memory **140** (e.g., the non-volatile memory arrays **205**) may include cache memory **225** (sometimes called non-volatile cache memory) and main memory **230** (sometimes called non-volatile main memory). The cache memory **225** and the main memory **230** may both be non-volatile memory. The cache memory **225** may store fewer bits per memory cell than the main memory **230**, and the main memory **230** may store more bits per memory cell than the cache memory **225**. For example, the cache memory **225** may include SLC memory cells that store one bit per cell, while the main memory **230** may include MLC memory cells that store more than one bit per cell, such as dual-level cell (DLC) memory cells that store two bits per cell, triple-level cell (TLC) memory cells that store three bits per cell, quad-level cell (QLC) memory cells that store four bits per cell, and/or penta-level cell (PLC) memory cells that store five bits per cell. Thus, the cache memory **225** may be referred to as SLC cache memory and the main memory **230** may be referred to as MLC main memory. In some implementations, the cache memory **225** (or a portion of the cache memory **225**) is used for write boosting. For example, the cache memory **225** (or a portion of the cache memory **225**) may be configured to store data for which write boosting is activated.

(29) As further shown in FIG. 2, the controller **130** may include a memory management component **235**, a memory reservation component **240**, a memory selection component **245**, and/or a command execution component **250**. In some implementations, one or more of these components are implemented as one or more instructions (e.g., firmware) executed by the controller **130**. Alternatively, one or more of these components may be implemented as dedicated integrated circuits distinct from the controller **130**.

(30) The memory management component **235** may be configured to manage performance of the memory device **120**. For example, the memory management component **235** may perform wear leveling, bad block management, block retirement, read disturb management, and/or other memory management operations. In some implementations, the memory device **120** may store (e.g., in memory **140**) one or more memory management tables. A memory management table may store information that may be used by or updated by the memory management component **235**, such as information regarding memory block age, memory block erase count, and/or error information associated with a memory partition (e.g., a memory cell, a row of memory, a block of memory, or the like).

(31) The memory reservation component **240** may be configured to reserve a portion of the cache memory **225** for data for which write boosting is activated (sometimes called “write-boosted data”). For example, the memory reservation component **240** may receive, from the host device **110**, an indication of an amount of cache memory **225** to be reserved for write-boosted data and may reserve that amount of the cache memory **225** for the write-boosted data.

(32) The memory selection component **245** may determine and/or select a portion of the memory **140** to which data is to be written. In some implementations, the memory selection component **245** may determine whether to write data to the cache memory **225** or the main memory **230** based on a data type of the data (e.g., hot data, cold data, or small fragment data). Additionally, or alternatively, the memory selection component **245** may determine whether to write data to the cache memory **225** or the main memory **230** based on whether the cache memory **225** has available memory that is not reserved for write-boosted data (e.g., whether writing data to the cache memory **225** would cause the available memory of the cache memory **225** to fall below the amount of cache memory **225** reserved for write-boosted data). Additionally, or alternatively, the memory selection component **245** may determine whether to write data to the cache memory **225** or the main memory **230** based on whether write boosting is activated for the data. In some implementations, the memory selection component **245** may determine to write data to the cache memory **225** despite write boosting being deactivated for the data (e.g., based on the data type and/or the available

memory of the cache memory **225**).

(33) The command execution component **250** may be configured to execute one or more memory commands, such as a read command, a write command (sometimes called a program command), or an erase command. For example, the command execution component **250** may receive and/or execute a write command to write data to the memory **140**. In some implementations, the command execution component **250** may receive an indication from the memory selection component **245** that indicates whether to write the data to the cache memory **225** or the main memory **230**, and the command execution component **250** may write the data to the cache memory **225** or the main memory **230** based on the indication. Additionally, or alternatively, the command execution component **250** may be configured to perform a garbage collection operation, such as by copying data from an old block (e.g., of cache memory **225**) to one or more new blocks (e.g., of cache memory **225** and/or main memory **230**).

(34) One or more devices or components shown in FIG. 2 may be configured to perform operations described elsewhere herein, such as one or more operations and/or methods described in connection with FIGS. 3-9. For example, the controller **130**, the memory management component **235**, the memory reservation component **240**, the memory selection component **245**, and/or the command execution component **250** may be configured to perform one or more operations and/or methods for the memory device **120**.

(35) The number and arrangement of components shown in FIG. 2 are provided as an example. In practice, there may be additional components, fewer components, different components, or differently arranged components than those shown in FIG. 2. Furthermore, two or more components shown in FIG. 2 may be implemented within a single component, or a single component shown in FIG. 2 may be implemented as multiple, distributed components. Additionally, or alternatively, a set of components (e.g., one or more components) shown in FIG. 2 may perform one or more operations described as being performed by another set of components shown in FIG. 2.

(36) FIG. 3 is a diagram illustrating an example memory architecture **300** that may be used by the memory device **120**. The memory device **120** may use the memory architecture **300** to store data. As shown, the memory architecture **300** may include a die **310**, which may include multiple planes **320**. A plane **320** may include multiple blocks **330**. A block **330** may include multiple pages **340**. Although FIG. 3 shows a particular quantity of planes **320** per die **310**, a particular quantity of blocks **330** per plane **320**, and a particular quantity of pages **340** per block **330**, these quantities may be different than what is shown. In some implementations, the memory architecture **300** is a NAND memory architecture.

(37) The die **310** is a structure made of semiconductor material, such as silicon. In some implementations, a die **310** is the smallest unit of memory that can independently execute commands. A memory device **120** may include one or more dies **310**. In some implementations, the memory device **120** may include multiple dies **310**. In this case, multiples dies **310** may each perform a respective memory operation (e.g., a read operation, a write operation, or an erase operation) in parallel. For example, a controller **130** of the memory device **120** may be configured to concurrently perform memory operations on multiple dies **310** for parallel control. In some implementations, a die **310** may include a memory array (e.g., a non-volatile memory array) and a local controller. Multiple dies **310** may be stacked and/or otherwise incorporated into a memory device **120** along with an external (or global) controller that communicates with respective local controllers of each die **310**.

(38) Each die **310** of a memory device **120** includes one or more planes **320**. A plane **320** is sometimes called a memory plane. In some implementations, identical and concurrent operations can be performed on multiple planes **320** (sometimes with restrictions). For example, a multi-plane command (e.g., a multi-plane read command or a multi-plane write command) may be executed on multiple planes **320** concurrently, whereas a single plane command (e.g., a single plane read

command or a single plane write command) may be executed on a single plane **320**. A logical unit of the memory device **120** may include one or more planes **320** of a die **310**. In some implementations, a logical unit may include all planes **320** of a die **310** and may be equivalent to a die **310**. Alternatively, a logical unit may include fewer than all planes **320** of a die **310**. A logical unit may be identified by a logical unit number (LUN). Depending on the context, the term “LUN” may refer to a logical unit or an identifier (e.g., a number) of that logical unit.

(39) Each plane **320** includes multiple blocks **330**. A block **330** is sometimes called a memory block. Each block **330** includes multiple pages **340**. A page **340** is sometimes called a memory page. A block **330** is the smallest unit of memory that can be erased. In other words, an individual page **340** of a block **330** cannot be erased without erasing every other page **340** of the block **330**. A page **340** is the smallest unit of memory to which data can be written (i.e., the smallest unit of memory that can be programmed with data). The terminology “programming” memory and “writing to” memory may be used interchangeably. A page **340** may include multiple memory cells that are accessible via the same access line (sometimes called a word line). A “page line” refers to a group of pages **340** at a same position across multiple planes **320** in a group of planes **320** (e.g., a first page having page index 0 in a first plane, a second page having page index 0 in a second plane, a third page having page index 0 in a third plane, and so on). In some implementations, a block **330** may be divided into multiple sub-blocks. A sub-block is a portion of a block **330** and may include a subset of pages **340** of the block and/or a subset of memory cells of the block **330**.

(40) In some implementations, read and write operations are performed for a specific page **340**, while erase operations are performed for a block **330** (e.g., all pages **340** in the block **330**). In some implementations, to prevent wearing out of memory, all pages **340** of a block **330** may be programmed before the block **330** is erased to enable a new program operation to be performed to a page **340** of the block **330**. After a page **340** is programmed with data (called “old data” below), that data can be erased, but that data cannot be overwritten with new data prior to being erased. The erase operation would erase all pages **340** in the block **330**, and erasing the entire block **330** every time that new data is to replace old data would quickly wear out the memory cells of the block **330**. Thus, rather than performing an erase operation, the new data may be stored in a new page (e.g., an empty page), as shown by reference number **350**, and the old page that stores the old data may be marked as invalid, as shown by reference number **360**. The memory device **120** may then point operations associated with the data to the new page (e.g., in an address table) and may track invalid pages to prevent program operations from being performed on invalid pages prior to an erase operation.

(41) When a block **330** satisfies an erasure condition, the memory device **120** may select the block **330** for erasure, copy the valid data of the block **330** (e.g., to a new block **330** or to the same block **330** after erasure), and erase the block **330**. For example, the erasure condition may be that all pages **340** of the block **330** or a threshold quantity or percentage of pages **340** of the block **330** are unavailable for further programming (e.g., are either invalid or already store valid data). As another example, the erasure condition may be that a quantity or percentage of free pages **340** of the block **330** (e.g., pages **340** that are available to be written) is less than or equal to a threshold. The process of selecting a block **330** satisfying an erasure condition, copying valid pages **340** of that block **330** to a new block **330** (or the same block **330** after erasure), and erasing the block **330** is sometimes called garbage collection and is used to free up memory space of the memory device **120**.

(42) As indicated above, FIG. 3 is provided as an example. Other examples may differ from what is described with regard to FIG. 3.

(43) FIG. 4 is a diagram of an example **400** of reserving a portion of write booster cache memory for write-boosted data. The operations described in connection with FIG. 4 may be performed by the memory device **120** and/or one or more components of the memory device **120**, such as the controller **130** and/or one or more components of the controller **130**.

(44) As shown by reference number **410**, the memory device **120** (e.g., the controller **130**) may

receive an indication of an amount of cache memory **225** (e.g., non-volatile cache memory) to be reserved for write boosting. In other words, the memory device **120** may receive an indication of an amount of cache memory **225** to be reserved for data (e.g., write-boosted data) for which write boosting is activated. As shown, in some implementations, the memory device **120** may receive the indication from the host device **110**. The amount of cache memory **225** to be reserved may be referred to as a threshold amount or a reserved amount. “Write boosting” is sometimes called “SLC data caching” or “write boost data caching” herein.

(45) In some implementations, the host device **110** may determine the amount of cache memory **225** to be reserved for write boosting based on one or more characteristics of the host device **110**, such as a device type of the host device **110** (e.g., a server, a personal computer, a smartphone, or an automotive device), a frequency with which the host device **110** is expected to request write boosting, and/or a rate at which the host device **110** is expected to transmit write commands to the memory device **120**. In some implementations, the threshold amount may be indicated as a threshold size of a portion of the cache memory **225** to be reserved, such as a quantity of bytes (e.g., 100 megabytes, 1 gigabyte, 5 gigabytes, or the like). Alternatively, the threshold amount may be indicated as a threshold percentage of the cache memory **225** to be reserved, such as ten percent, twenty percent, thirty percent, or the like. In some implementations, the host device **110** may indicate the amount of cache memory **225** to be reserved by indicating the amount of cache memory **225** to be made available for non-write-boosted data, with the remaining amount of cache memory **225** to be reserved for write-boosted data.

(46) Additionally, or alternatively, the host device **110** may indicate one or more characteristics (as described above) of the host device **110** to the memory device **120**, and the memory device **120** may determine the amount of cache memory **225** to be reserved for write boosting based on the one or more characteristics. For example, the memory device **120** may reserve a larger quantity or percentage of the cache memory **225** for a host device **110** that is expected to request write boosting more frequently, and the memory device **120** may reserve a smaller quantity or percentage of the cache memory **225** for a host device **110** that is expected to request write boosting less frequently.

(47) As shown by reference number **420**, the memory device **120** (e.g., the controller **130**) may configure the cache memory **225** based on the indication, such as by reserving the threshold amount of cache memory **225** for write boosting. For example, the memory device **120** may store an indication, in the memory **140**, of the reserved amount of cache memory **225** reserved for write-boosted data. Additionally, or alternatively, the memory device **120** may store an indication, in the memory **140**, of an unreserved amount of cache memory **225** that is not reserved for write-boosted data. In the example **400** of FIG. **4**, the memory device **120** reserves twenty percent of the cache memory **225** for write-boosted data (shown as “WB data”), and the remaining eighty percent of the cache memory **225** is unreserved and available for non-write-boosted data.

(48) As used herein, “write-boosted data” refers to first data for which write boosting is activated, and “non-write-boosted data” refers to second data for which write boosting is deactivated. Write boosting may be activated or deactivated for data by the host device **110**. For example, the host device **110** may activate write boosting for first data by transmitting an activation instruction to the memory device **120** (e.g., in a write command or prior to transmitting the write command) that instructs the memory device **120** to activate write boosting for a write command that includes the first data. In this case, any write command not indicated by the activation instruction is associated with second data for which write boosting is deactivated. As another example, the host device **110** may transmit an activation instruction to activate write boosting, may transmit one or more write commands for which write boosting is activated, and may then transmit a deactivation instruction to deactivate write boosting (e.g., for any subsequent write commands until another activation instruction is transmitted).

(49) As shown by reference number **430**, in some implementations, the memory device **120** (e.g.,

the controller **130**) may reserve a portion of the cache memory **225** for small fragment data. For example, and as shown, the memory device **120** may reserve a portion of the unreserved cache memory **225**, that is not reserved for write-boosted data, for small fragment data. Small fragment data is a type of data having a size (e.g., in bytes) that is less than a threshold size. For example, small fragment data may have a size that is less than a page line size of a page line of the memory device **120** (e.g., a page line of the main memory **230**). A page line may include a group of pages (e.g., of the main memory **230**) located at a same position and/or identified by a same index value across multiple planes and/or multiple dies of the memory device **120**. For example, if a page size (e.g., an amount of data stored in an individual page) is 16 kilobytes and there are 16 pages in a page line, then the page line size is 256 kilobytes. In this case, small fragment data may be data, to be written to the memory device **120**, having a size less than 256 kilobytes.

(50) In some cases, the memory device **120** may be configured to write data to the main memory **230** such that an entire page line is filled with each write operation. If the memory device **120** receives a write command indicating small fragment data (e.g., less than 256 kilobytes) to be written to the memory device **120**, then the memory device **120** may need to write non-host data (e.g., dummy data) to the remaining pages of the page line that are not occupied by the small fragment data. This wastes memory resources of the main memory **230**. To improve memory utilization, the memory device **120** may store the small fragment data in the cache memory **225** until a condition is satisfied (e.g., enough small fragment data has accumulated in the cache memory **225** to fill an entire page line or a threshold portion of a page line).

(51) By reserving a first portion of the cache memory **225** for write-boosted data while enabling use of a second portion of the cache memory **225** for non-write-boosted data (particularly hot data, as described in more detail below), the memory device **120** may extend the lifespan of the memory device **120** (e.g., by reducing write amplification of hot data in the main memory **230**), may improve resource utilization of the memory device **120** (e.g., by permitting flexible use of the cache memory **225** that would otherwise be entirely used for write boosting), and/or may improve performance of the memory device **120** (e.g., by ensuring that a portion of the cache memory **225** is available for write-boosting operations). Furthermore, by reserving a portion of the cache memory **225** for small fragment data, the memory device **120** may ensure that the small fragment data operations described above can be performed, thereby improving memory resource utilization of the main memory **230**.

(52) As indicated above, FIG. 4 is provided as an example. Other examples may differ from what is described with regard to FIG. 4.

(53) FIG. 5 is a diagram of an example **500** of opportunistic storage of non-write-boosted data in write booster cache memory. The operations described in connection with FIG. 5 may be performed by the memory device **120** and/or one or more components of the memory device **120**, such as the controller **130** and/or one or more components of the controller **130**.

(54) As shown by reference number **510**, the memory device **120** (e.g., the controller **130**) may receive a write command for non-write-boosted data (e.g., data for which write boosting is deactivated). As shown, in some implementations, the memory device **120** may receive the write command from the host device **110**. In some implementations, the write command includes the data and an indication that write boosting is deactivated for the data. In some implementations, the write command does not include an indication that write boosting is deactivated for the data. Rather, the host device **110** may transmit a separate instruction that indicates whether write boosting is activated for a write command. In some implementations, write boosting may be deactivated by default, and the memory device **120** may determine that write boosting is deactivated for a write command unless the memory device **120** receives an indication that write boosting is activated for the write command.

(55) As shown by reference number **520**, the memory device **120** (e.g., the controller **130**) may determine a data type of the non-write-boosted data. The data type may be a first data type, such as

hot data, or a second data type, such as cold data. As used herein, “hot data” may include data with a relatively short expected lifespan (e.g., before being overwritten or marked as invalid, as described above in connection with FIG. 3), and “cold data” may include data with a relatively long expected lifespan. For example, hot data may have an expected lifespan that is shorter than an expected lifespan associated with cold data. In other words, hot data is expected to become invalid and/or is expected to be overwritten in a shorter amount of time than cold data. In some implementations, hot data has an expected lifespan that is less than (or equal to) a threshold, and cold data has an expected lifespan that is greater than (or equal to) the threshold.

(56) Additionally, or alternatively, “hot data” may include data that is accessed, erased, and/or overwritten more frequently (e.g., than cold data) or relatively frequently (e.g., compared to cold data), and “cold data” may include data that is accessed, erased, and/or overwritten less frequently (e.g., than hot data) or relatively infrequently (e.g., compared to hot data). For example, because hot data is erased more frequently than cold data, hot data may have a shorter lifespan (e.g., in a particular memory location) than cold data because the hot data stored in a particular memory location may be erased or marked as invalid, and new data that replaces the hot data stored in the particular memory location may be written to a new memory location. In some implementations, hot data is accessed, erased, and/or overwritten with a frequency (e.g., a quantity of instances in a time period) that is greater than (or equal to) a threshold, and cold data is accessed, erased, and/or overwritten with a frequency that is less than (or equal to) a threshold.

(57) In some implementations, the memory device **120** may receive an indication of whether non-write-boosted data included in the write command (or associated with the write command) is hot data or cold data. For example, the host device **110** may indicate (e.g., in the write command or in a separate message) whether the non-write-boosted data included in the write command is hot data or cold data. Additionally, or alternatively, the memory device **120** may determine whether the non-write-boosted data included in the write command is hot data or cold data based on one or more characteristics of the non-write-boosted data (e.g., a size of the non-write-boosted data), one or more characteristics of the write command (e.g., other information included in the write command and/or a rate at which write commands are received from the host device **110**), and/or one or more characteristics associated with the host device **110** (as described elsewhere herein).

(58) In some implementations, the memory device **120** may use historical information to determine whether the non-write-boosted data in a write command is hot data or cold data. For example, the memory device **120** may store information about historical data written by the memory device **120** based on historical write commands received from the host device **110**. The memory device **120** may store, for example, an indication of an actual lifespan of historical data previously written to the memory device **120** (e.g., a time between writing the data and overwriting or marking the data as invalid) along with one or more historical characteristics of the data, one or more historical characteristics of the write command in which the data was received, and/or one or more historical characteristics of the host device **110**. The memory device **120** may compare one or more characteristics associated with a new write command (that includes non-write-boosted data) and one or more historical characteristics (and corresponding actual lifespans) to determine whether the non-write-boosted data is hot data or cold data.

(59) Additionally, or alternatively, the data type may be small fragment data, as described above in connection with FIG. 4. In some implementations, the memory device **120** may determine whether the non-write-boosted data is small fragment data based on comparing the size of the non-write-boosted data to a threshold size, such as a page line size. If the size of the non-write-boosted data is less than (or equal to) the threshold size, then the non-write-boosted data is small fragment data. If the size of the non-write-boosted data is greater than (or equal to) the threshold size, then the non-write-boosted data is not small fragment data. Although data cannot be both hot data and cold data, in some implementations, data may be both hot data and small fragment data (e.g., small fragment data with a relatively short expected lifespan) or may be both cold data and small fragment data

(e.g., small fragment data with a relatively long expected lifespan).

(60) As shown by reference number **530**, the memory device **120** may determine whether there is sufficient unreserved cache memory **225** available for storage of the non-write-boostered data. As described above in connection with FIG. **4**, the memory device **120** may reserve a threshold amount of the cache memory **225** for write-boostered data. If the memory device **120** determines that writing the non-write-boostered data to the cache memory **225** would cause the available memory of the cache memory **225** to fall below the threshold amount, then the memory device **120** may refrain from writing the non-write-boostered data to the cache memory **225**. In this case, the memory device **120** may write the non-write-boostered data to the main memory **230**. Additionally, or alternatively, if the memory device **120** determines that writing the non-write-boostered data to the cache memory **225** would leave a remaining amount of available memory, of the cache memory **225**, that satisfies a threshold (e.g., that is greater than or equal to the threshold amount), then the memory device **120** may write the non-write-boostered data to the cache memory **225** (e.g., if the non-write-boostered data is hot data and/or small fragment data, as described elsewhere herein).

(61) As another example, the memory device **120** may determine a threshold amount of unreserved cache memory **225** (e.g., that is not reserved for write-boostered data). If the memory device **120** determines that writing the non-write-boostered data to the cache memory **225** would cause the amount of the cache memory **225** used for non-write-boostered data to exceed this threshold amount, then the memory device **120** may refrain from writing the non-write-boostered data to the cache memory **225** (and may write the non-write-boostered data to the main memory **230**). In some implementations, the threshold amount of unreserved cache memory **225** available for non-write-boostered data is based on a threshold amount of cache memory **225** reserved for write-boostered data (e.g., a remaining amount of unreserved cache memory **225**). In some implementations, the threshold amount of unreserved cache memory **225** available for non-write-boostered data is based on an amount of cache memory **225** reserved for write-boostered data and an amount of cache memory **225** reserved for small fragment data (e.g., a remaining amount of unreserved cache memory **225**).

(62) If the memory device **120** reserves a portion of the cache memory **225** for small fragment data and if the memory device **120** determines that the non-write-boostered data is small fragment data, then the memory device **120** may refrain from determining if the non-write-boostered data is hot data or cold data and may refrain from determining whether there is sufficient unreserved cache memory **225** available for non-write-boostered data. This may result in faster write times by removing unnecessary operations. However, in some implementations, the memory device **120** may determine whether there is sufficient cache memory **225** available for the small fragment data (e.g., using the cache memory **225** reserved for the small fragment data and/or the unreserved cache memory **225**). If there is not, then the memory device **120** may write the small fragment data to the main memory **230**.

(63) In some implementations, the memory device **120** may determine whether there is sufficient unreserved cache memory **225** available for storage of the non-write-boostered data (as described in connection with reference number **530**) before determining whether the non-write-boostered data is hot data or cold data (as described in connection with reference number **520**). If there is insufficient unreserved cache memory **225** available for storage of the non-write-boostered data, then the memory device **120** may refrain from determining whether the non-write-boostered data is hot data or cold data. This may result in faster write times by removing unnecessary operations because the memory device **120** cannot write the non-write-boostered data to the cache memory **225** regardless of whether the non-write-boostered data is hot data or cold data. However, in some implementations, the memory device **120** may still determine whether the non-write-boostered data is small fragment data, as described above. Furthermore, if the memory device **120** determines that there is sufficient unreserved cache memory **225** available for storage of the non-write-boostered data, the memory device **120** may then determine whether the non-write-boostered data is hot data or cold data (as

described in connection with reference number 520).

(64) Alternatively, the memory device **120** may determine whether the non-write-boosted data is hot data or cold data before determining whether there is sufficient unreserved cache memory **225** available for storage of the non-write-boosted data. If the non-write-boosted data is cold data, then the memory device **120** may refrain from determining whether there is sufficient unreserved cache memory **225** available for storage of the non-write-boosted data. This may result in faster write times by removing unnecessary operations because the memory device **120** does not write cold non-write-boosted data to the cache memory **225**. However, if the memory device **120** determines that the non-write-boosted data is hot data, the memory device **120** may then determine whether there is sufficient unreserved cache memory **225** available for storage of the hot non-write-boosted data.

(65) As shown by reference number **540**, the memory device **120** may selectively write the non-write-boosted data to the cache memory **225** or the main memory **230** based on the data type of the non-write boosted data and/or based on whether sufficient unreserved cache memory **225** is available. As used herein, “selectively” writing the non-write-boosted data to the cache memory **225** or the main memory **230** means to either write the non-write-boosted data to the cache memory **225** or to write the non-write-boosted data to the main memory **230**. For example, selectively writing the non-write-boosted data to the cache memory **225** or the main memory **230** based on whether a condition is satisfied (e.g., a condition dependent on the data type and/or whether sufficient unreserved cache memory **225** is available) means that the non-write-boosted data is written to the cache memory **225** if the condition is satisfied and that the non-write-boosted data is written to the main memory **230** if the condition is not satisfied (or vice versa). Thus, selectively writing the non-write-boosted data to the cache memory **225** or the main memory **230** may include determining whether to write the non-write-boosted data to the cache memory **225** or whether to write the non-write-boosted data to the main memory **230** and then either writing the non-write-boosted data to the cache memory **225** or writing the non-write-boosted data to the main memory **230** based on that determination.

(66) As an example, and as shown by reference number **550**, if the memory device **120** determines that the non-write-boosted data is hot data (e.g., is associated with a first data type) and that the cache memory **225** has sufficient memory (e.g., that is not reserved for write-boosted data and/or that is not reserved for small fragment data) to store the hot non-write-boosted data, then the memory device **120** may write the hot non-write-boosted data to the cache memory **225** (e.g., to store the hot non-write-boosted data in the cache memory **225**). Additionally, or alternatively, if the memory device **120** determines that the non-write-boosted data is small fragment data and that the cache memory **225** has sufficient memory (e.g., that is not reserved for write-boosted data and/or that is reserved for small fragment data) to store the small fragment non-write-boosted data, then the memory device **120** may write the small fragment non-write-boosted data to the cache memory **225** (e.g., to store the small fragment non-write-boosted data in the cache memory **225**). The memory device **120** may write the non-write-boosted data to the cache memory **225** (e.g., the write booster cache memory) despite write boosting being deactivated for the non-write-boosted data.

(67) In some implementations, the memory device **120** may use a write boost cursor to identify a location of the cache memory **225** to which the non-write-boosted data is to be written. A “write boost cursor” may include information stored by the memory device **120** (e.g., in non-volatile memory) that indicates the location. The location may include, for example, a die, a plane, a block, a sub-block, a page line, and/or a page. If the memory device **120** determines that the non-write-boosted data is to be stored in the cache memory **225**, then the memory device **120** may use the write boost cursor (sometimes called a cache memory cursor) to determine the location to which the non-write boosted data is to be written. Otherwise, the memory device **120** may use a different write cursor (e.g., a main memory cursor) to determine a memory location to which the non-write-boosted data is to be written (e.g., in the main memory **230**).

(68) As another example, and as shown by reference number **560**, if the memory device **120** determines that the non-write-boostered data is cold data (e.g., is associated with a second data type) or that the cache memory **225** does not have sufficient memory (e.g., that is not reserved for write-boostered data and/or that is not reserved for small fragment data) to store the non-write-boostered data, then the memory device **120** may write the non-write-boostered data to the main memory **230** (e.g., to store the non-write-boostered data in the main memory **230**). As described above, if the memory device **120** determines that the non-write-boostered data is to be stored in the main memory **230**, then the memory device **120** may use a main memory cursor to determine the location (e.g., a die, a plane, a block, a sub-block, a page line, and/or a page) in the main memory **230** to which the non-write boostered data is to be written.

(69) By writing some non-write-boostered data (e.g., hot data and/or small fragment data) to the cache memory **225** while also reserving a portion of the cache memory **225** for write-boostered data, the memory device **120** may extend the lifespan of the memory device **120** (e.g., by reducing write amplification of hot data in the main memory **230**), may improve resource utilization of the memory device **120** (e.g., by permitting flexible use of the cache memory **225** that would otherwise be entirely used for write boosting and by refraining from writing large amounts of dummy data along with small fragment data to the main memory **230**), and/or may improve performance of the memory device **120** (e.g., by ensuring that a portion of the cache memory **225** is available for write-boosting operations).

(70) As indicated above, FIG. 5 is provided as an example. Other examples may differ from what is described with regard to FIG. 5.

(71) FIG. 6 is a diagram of an example **600** of performing a garbage collection operation on write booster cache memory. The operations described in connection with FIG. 6 may be performed by the memory device **120** and/or one or more components of the memory device **120**, such as the controller **130** and/or one or more components of the controller **130**.

(72) As shown by reference number **610**, the memory device **120** (e.g., the controller **130**) may perform a garbage collection operation on the cache memory **225**. As described above in connection with FIG. 3, the garbage collection operation may include, for example, identifying a block (sometimes called an old block) of cache memory **225** that satisfies an erasure condition (as described above in connection with FIG. 3), copying valid pages of the identified block to one or more new blocks, and erasing the identified block. The memory device **120** may perform the garbage collection to free up memory resources of the cache memory **225**.

(73) In some implementations, as part of performing the garbage collection operation, the memory device **120** may write data, from the old block, having a first data type (e.g., hot data) to a first block **620** and may write data, from the old block, having a second data type (e.g., cold data) to a second block **630**. This may reduce write amplification and wear-out of blocks because copying all of the hot data to a single block, rather than multiple blocks, may reduce a quantity of blocks that need to then be later erased within a short time period of copying the data.

(74) For example, as shown by reference number **640**, the memory device **120** may copy hot data from the old block of cache memory **225** to the first block **620**. To identify the hot data from the old block, the memory device **120** may identify a first set of pages, included in the old block, that store hot data (e.g., first data associated with a first data type). The memory device **120** may copy the hot data to the first block **620** based on identifying the first set of pages.

(75) Similarly, as shown by reference number **650**, the memory device **120** may copy cold data from the old block of cache memory **225** to the second block **630**. To identify the cold data from the old block, the memory device **120** may identify a second set of pages, included in the old block, that store cold data (e.g., second data associated with a second data type). The memory device **120** may copy the cold data to the second block **630** based on identifying the second set of pages.

(76) In some implementations, the memory device **120** may store a data type indicator that indicates, for one or more pages of the cache memory **225** (e.g., an unreserved portion of the cache

memory **225**), whether data stored in those one or more pages is hot data or cold data. For example, the data type indicator may have a first value (e.g., 0) when corresponding data is hot data, and the data type indicator may have a second value (e.g., 1) when corresponding data is cold data. Additionally, or alternatively, the data type indicator may indicate a time at which data was written (e.g., to one or more pages), and the memory device **120** may use that time to determine whether the data is hot data or cold data. For example, the memory device **120** may compare the time (e.g., a stored time) to a current time (e.g., associated with performing the garbage collection operation) to determine whether data stored in a corresponding page is cold data (e.g., with a difference between the stored time and current time that satisfies a threshold) or hot data (e.g., with a difference between the stored time and current time that does not satisfy the threshold). The memory device **120** may use one or more data type indicators to identify the first set of pages that store the hot data and/or the second set of pages that store the cold data.

(77) In some implementations, the first block **620** and the second block **630** are both included in the main memory **230**, and both the hot data and the cold data are copied to the main memory **230**. This may free up more memory resources of the cache memory **225** than if some of the data (e.g., the hot data) is stored in a new block of the cache memory **225**. Alternatively, in some implementations, the first block **620** may be included in the cache memory **225** and the second block **630** may be included in the main memory **230**. This may free up some of the cache memory **225** (e.g., by copying cold data to the main memory **230**) while reducing write amplification of the main memory **230** (e.g., by keeping the hot data in the cache memory **225**).

(78) In some implementations, the memory device **120** may determine whether to copy both the hot data and the cold data to the main memory **230** or to copy the cold data to the main memory **230** and the hot data to the cache memory **225** based on an amount of available unreserved cache memory **225**. For example, if the amount of available unreserved cache memory **225** is less than (or equal to) a threshold amount, then the memory device **120** may copy both the hot data and the cold data to the main memory **230** to free up memory resources of the cache memory **225** (e.g., to create more available unreserved cache memory **225**). If the amount of available unreserved cache memory **225** is greater than (or equal to) the threshold amount, then the memory device **120** may copy the cold data to the main memory **230** and may copy the hot data to the cache memory **225** to reduce write amplification of the main memory **230** (and because there is less need to free up memory resources of the cache memory **225**).

(79) Additionally, or alternatively, the memory device **120** may determine whether to copy both the hot data and the cold data to the main memory **230** or to copy the cold data to the main memory **230** and the hot data to the cache memory **225** based on one or more block age parameters associated with the cache memory **225**, such as a number of P/E cycles associated with available blocks of the cache memory **225** (e.g., an unreserved portion of the cache memory **225**), and an average number of P/E cycles associated with blocks of the cache memory **225** (e.g., an unreserved portion of the cache memory **225**). For example, if a P/E cycle parameter associated with the cache memory **225** is greater than (or equal to) a threshold, then the memory device **120** may copy both the hot data and the cold data to the main memory **230** to reduce wear on the cache memory **225**. If the P/E cycle parameter associated with the cache memory **225** is less than (or equal to) the threshold, then the memory device **120** may copy the cold data to the main memory **230** and may copy the hot data to the cache memory **225** to reduce write amplification of the main memory **230** (and because there is less need to reduce wear on the cache memory **225**).

(80) In some implementations, the memory device **120** may use one or more block age parameters to select a block of the cache memory **225** and/or a block of the main memory **230** to which to copy hot data. For example, the memory device **120** may copy hot data to a block with a lower P/E cycle count (e.g., less than or equal to a threshold) and/or may refrain from copying hot data to a block with a higher P/E cycle count (e.g., greater than or equal to a threshold). This may result in more even wear on blocks of the cache memory **225** and/or the main memory **230**.

(81) In some implementations, the memory device **120** may modify a configuration that indicates a first set of blocks of the cache memory **225** reserved for write-boosted data, a second set of blocks of the cache memory **225** reserved for small fragment data, and/or a third set of blocks of the cache memory **225** that are unreserved. For example, if one of the sets of blocks is wearing out faster than the other set(s) of blocks (e.g., based on a P/E cycle count), then the memory device **120** may modify which blocks are included in each set (e.g., to achieve wear leveling).

(82) As indicated above, FIG. **6** is provided as an example. Other examples may differ from what is described with regard to FIG. **6**.

(83) FIG. **7** is a flowchart of an example method **700** associated with opportunistic storage of non-write-boosted data in write booster cache memory. In some implementations, a memory device (e.g., the memory device **120**) may perform or may be configured to perform the method **700**. In some implementations, another device or a group of devices separate from or including the memory device (e.g., the system **100**) may perform or may be configured to perform the method **700**. Additionally, or alternatively, one or more components of the memory device (e.g., the controller **130**, the memory management component **235**, the memory reservation component **240**, the memory selection component **245**, and/or the command execution component **250**) may perform or may be configured to perform the method **700**. Thus, means for performing the method **700** may include the memory device and/or one or more components of the memory device. Additionally, or alternatively, a non-transitory computer-readable medium may store one or more instructions that, when executed by the memory device (e.g., the controller **130** of the memory device **120**), cause the memory device to perform the method **700**.

(84) As shown in FIG. **7**, the method **700** may include receiving, from a host device, an indication of a threshold amount of non-volatile cache memory to be reserved for first data for which write boosting is activated (block **710**). As further shown in FIG. **7**, the method **700** may include receiving, from the host device, a write command that includes second data for which write boosting is deactivated (block **720**). As further shown in FIG. **7**, the method **700** may include determining that the second data is hot data (block **730**). As further shown in FIG. **7**, the method **700** may include identifying a block of the non-volatile cache memory using a write boost cursor based on determining that the second data is hot data, based on the indication of the threshold amount, and despite write boosting being deactivated for the second data (block **740**). As further shown in FIG. **7**, the method **700** may include writing the second data to the block of the non-volatile cache memory (block **750**).

(85) The method **700** may include additional aspects, such as any single aspect or any combination of aspects described below and/or described in connection with one or more other methods or operations described elsewhere herein.

(86) In a first aspect, the method **700** includes identifying the block using the write boost cursor based on a determination that writing the second data to the non-volatile cache memory will not reduce an available amount of non-volatile cache memory reserved for the first data below the threshold amount.

(87) In a second aspect, alone or in combination with the first aspect, the method **700** includes copying hot data stored in the block to a first new block and copying cold data stored in the block to a second new block during a garbage collection operation performed on the block.

(88) In a third aspect, alone or in combination with one or more of the first and second aspects, the first new block is a first block of non-volatile main memory and the second new block is a second block of the non-volatile main memory.

(89) In a fourth aspect, alone or in combination with one or more of the first through third aspects, the first new block is a new block of the non-volatile cache memory and the second new block is a new block of non-volatile main memory.

(90) In a fifth aspect, alone or in combination with one or more of the first through fourth aspects, the method **700** includes copying the hot data to the first new block and copying the cold data to

the second new block during the garbage collection operation based on one or more block age parameters associated with the non-volatile cache memory.

(91) In a sixth aspect, alone or in combination with one or more of the first through fifth aspects, the indication of the threshold amount comprises a threshold size or a threshold percentage.

(92) In a seventh aspect, alone or in combination with one or more of the first through sixth aspects, the method **700** includes storing small fragment data, having a size less than a page line size of a page line of the memory device, in the block of the non-volatile cache memory.

(93) In an eighth aspect, alone or in combination with one or more of the first through seventh aspects, a portion of the non-volatile cache memory is reserved for the small fragment data.

(94) Although FIG. 7 shows example blocks of a method **700**, in some implementations, the method **700** may include additional blocks, fewer blocks, different blocks, or differently arranged blocks than those depicted in FIG. 7. Additionally, or alternatively, two or more of the blocks of the method **700** may be performed in parallel. The method **700** is an example of one method that may be performed by one or more devices described herein. These one or more devices may perform or may be configured to perform one or more other methods based on operations described herein.

(95) FIG. 8 is a flowchart of an example method **800** associated with opportunistic storage of non-write-boosted data in write booster cache memory. In some implementations, a memory device (e.g., the memory device **120**) may perform or may be configured to perform the method **800**. In some implementations, another device or a group of devices separate from or including the memory device (e.g., the system **100**) may perform or may be configured to perform the method **800**. Additionally, or alternatively, one or more components of the memory device (e.g., the controller **130**, the memory management component **235**, the memory reservation component **240**, the memory selection component **245**, and/or the command execution component **250**) may perform or may be configured to perform the method **800**. Thus, means for performing the method **800** may include the memory device and/or one or more components of the memory device.

Additionally, or alternatively, a non-transitory computer-readable medium may store one or more instructions that, when executed by the memory device (e.g., the controller **130** of the memory device **120**), cause the memory device to perform the method **800**.

(96) As shown in FIG. 8, the method **800** may include receiving a write command that includes data to be written to the memory device (block **810**). As further shown in FIG. 8, the method **800** may include receiving an indication that single-level cell data caching is deactivated for the data (block **820**). As further shown in FIG. 8, the method **800** may include determining whether the data is associated with a first data type or a second data type (block **830**). As further shown in FIG. 8, the method **800** may include selectively writing the data to single-level cell cache memory or multi-level cell main memory based on: a determination of whether the data is associated with the first data type or the second data type, and a determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching (block **840**).

(97) The method **800** may include additional aspects, such as any single aspect or any combination of aspects described below and/or described in connection with one or more other methods or operations described elsewhere herein.

(98) In a first aspect, the first data type is associated with an expected lifespan that is shorter than an expected lifespan associated with the second data type.

(99) In a second aspect, alone or in combination with the first aspect, selectively writing the data to the single-level cell cache memory or the multi-level cell main memory comprises writing the data to the single-level cell cache memory based on a determination that the data is associated with the first data type, and a determination that the single-level cell cache memory has sufficient memory, that is not reserved for the single-level cell data caching, to store the data.

(100) In a third aspect, alone or in combination with one or more of the first and second aspects, selectively writing the data to the single-level cell cache memory or the multi-level cell main

memory comprises writing the data to the multi-level cell main memory based on a determination that the data is associated with the second data type.

(101) In a fourth aspect, alone or in combination with one or more of the first through third aspects, selectively writing the data to the single-level cell cache memory or the multi-level cell main memory comprises writing the data to the multi-level cell main memory based on a determination that the single-level cell cache memory does not have sufficient memory, that is not reserved for the single-level cell data caching, to store the data.

(102) In a fifth aspect, alone or in combination with one or more of the first through fourth aspects, the method **800** includes receiving, from a host device, an indication of an amount of memory in the single-level cell cache memory to be reserved for the single-level cell data caching, and the determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching is based on the amount of memory in the single-level cell cache memory to be reserved for the single-level cell data caching.

(103) In a sixth aspect, alone or in combination with one or more of the first through fifth aspects, the method **800** includes identifying a block of the single-level cell cache memory for which a garbage collection operation is to be performed, identifying a first set of pages, included in the block, that store first data associated with the first data type, identifying a second set of pages, included in the block, that store second data associated with the second data type, copying the first data to a first block as part of performing the garbage collection operation, and copying the second data to a second block as part of performing the garbage collection operation.

(104) In a seventh aspect, alone or in combination with one or more of the first through sixth aspects, both the first block and the second block are included in the multi-level cell main memory.

(105) In an eighth aspect, alone or in combination with one or more of the first through seventh aspects, the first block is included in the single-level cell cache memory and the second block is included in the multi-level cell main memory.

(106) Although FIG. **8** shows example blocks of a method **800**, in some implementations, the method **800** may include additional blocks, fewer blocks, different blocks, or differently arranged blocks than those depicted in FIG. **8**. Additionally, or alternatively, two or more of the blocks of the method **800** may be performed in parallel. The method **800** is an example of one method that may be performed by one or more devices described herein. These one or more devices may perform or may be configured to perform one or more other methods based on operations described herein.

(107) FIG. **9** is a flowchart of an example method **900** associated with opportunistic storage of non-write-boosted data in write booster cache memory. In some implementations, a memory device (e.g., the memory device **120**) may perform or may be configured to perform the method **900**. In some implementations, another device or a group of devices separate from or including the memory device (e.g., the system **100**) may perform or may be configured to perform the method **900**. Additionally, or alternatively, one or more components of the memory device (e.g., the controller **130**, the memory management component **235**, the memory reservation component **240**, the memory selection component **245**, and/or the command execution component **250**) may perform or may be configured to perform the method **900**. Thus, means for performing the method **900** may include the memory device and/or one or more components of the memory device. Additionally, or alternatively, a non-transitory computer-readable medium may store one or more instructions that, when executed by the memory device (e.g., the controller **130** of the memory device **120**), cause the memory device to perform the method **900**.

(108) As shown in FIG. **9**, the method **900** may include receiving a write command that includes data to be written to the memory device (block **910**). As further shown in FIG. **9**, the method **900** may include determining that write boost data caching is deactivated for the data (block **920**). As further shown in FIG. **9**, the method **900** may include determining that the data is associated with a first data type (block **930**). As further shown in FIG. **9**, the method **900** may include determining that the data is to be written to cache memory, rather than main memory, based on determining that

the data is associated with the first data type and despite write boost data caching being deactivated for the data (block **940**). As further shown in FIG. **9**, the method **900** may include writing the data to the cache memory based on determining that the data is to be written to the cache memory (block **950**).

(109) The method **900** may include additional aspects, such as any single aspect or any combination of aspects described below and/or described in connection with one or more other methods or operations described elsewhere herein.

(110) In a first aspect, the first data type is expected to become invalid or is expected to be overwritten in a shorter amount of time than a second data type.

(111) In a second aspect, alone or in combination with the first aspect, the main memory stores more bits per memory cell than the cache memory.

(112) In a third aspect, alone or in combination with one or more of the first and second aspects, determining that the data is to be written to the cache memory further comprises determining that the data is to be written to the cache memory based on a determination that writing the data to the cache memory would leave a remaining amount of available cache memory that satisfies a threshold.

(113) In a fourth aspect, alone or in combination with one or more of the first through third aspects, the method **900** includes receiving an indication of the threshold, wherein the threshold indicates an amount of cache memory to be reserved for the write boost data caching.

(114) In a fifth aspect, alone or in combination with one or more of the first through fourth aspects, the method **900** includes writing small fragment data, having a size less than a page line size of a page line of the memory device, to the cache memory.

(115) In a sixth aspect, alone or in combination with one or more of the first through fifth aspects, the method **900** includes identifying a block of the cache memory for which a garbage collection operation is to be performed, and copying first data, associated with the first data type and stored in the block, to a first block as part of performing the garbage collection operation, and copying second data, associated with a first data type and stored in the block, to a second block as part of performing the garbage collection operation.

(116) Although FIG. **9** shows example blocks of a method **900**, in some implementations, the method **900** may include additional blocks, fewer blocks, different blocks, or differently arranged blocks than those depicted in FIG. **9**. Additionally, or alternatively, two or more of the blocks of the method **900** may be performed in parallel. The method **900** is an example of one method that may be performed by one or more devices described herein. These one or more devices may perform or may be configured to perform one or more other methods based on operations described herein.

(117) In some implementations, a memory device includes one or more components configured to: receive, from a host device, an indication of a threshold amount of non-volatile cache memory to be reserved for first data for which write boosting is activated; receive, from the host device, a write command that includes second data for which write boosting is deactivated; determine that the second data is hot data; identify a block of the non-volatile cache memory using a write boost cursor based on determining that the second data is hot data, based on the indication of the threshold amount, and despite write boosting being deactivated for the second data; and write the second data to the block of the non-volatile cache memory.

(118) In some implementations, a method includes receiving, by a memory device, a write command that includes data to be written to the memory device; receiving, by the memory device, an indication that single-level cell data caching is deactivated for the data; determining, by the memory device, whether the data is associated with a first data type or a second data type; and selectively writing, by the memory device, the data to single-level cell cache memory or multi-level cell main memory based on: a determination of whether the data is associated with the first data type or the second data type, and a determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching.

(119) In some implementations, an apparatus includes a non-volatile memory array; and a controller configured to execute instructions that cause the apparatus to: process a received command to write data to the non-volatile memory array, wherein the command is received without an indication that the data is to be written in a single-level cell (SLC) cache memory of the non-volatile memory array; determine whether the data is associated with a first data type; determine whether the SLC cache memory has sufficient memory not reserved for SLC data caching; and write the data to the SLC cache memory based on a determination that the data is associated with the first data type and that the SLC cache memory has sufficient memory not reserved for SLC data caching.

(120) The foregoing disclosure provides illustration and description but is not intended to be exhaustive or to limit the implementations to the precise forms disclosed. Modifications and variations may be made in light of the above disclosure or may be acquired from practice of the implementations described herein.

(121) As used herein, “satisfying a threshold” may, depending on the context, refer to a value being greater than the threshold, greater than or equal to the threshold, less than the threshold, less than or equal to the threshold, equal to the threshold, not equal to the threshold, or the like. The phrases “based on” or “in response to” may be used in place of “if” as used herein. For example, an operation being performed “if” a condition is satisfied may be replaced with the operation being performed “based on” the condition being satisfied or the operation being performed “in response to” the condition being satisfied.

(122) Even though particular combinations of features are recited in the claims and/or disclosed in the specification, these combinations are not intended to limit the disclosure of implementations described herein. Many of these features may be combined in ways not specifically recited in the claims and/or disclosed in the specification. For example, the disclosure includes each dependent claim in a claim set in combination with every other individual claim in that claim set and every combination of multiple claims in that claim set. As used herein, a phrase referring to “at least one of” a list of items refers to any combination of those items, including single members. As an example, “at least one of: a, b, or c” is intended to cover a, b, c, a+b, a+c, b+c, and a+b+c, as well as any combination with multiples of the same element (e.g., a+a, a+a+a, a+a+b, a+a+c, a+b+b, a+c+c, b+b, b+b+b, b+b+c, c+c, and c+c+c, or any other ordering of a, b, and c).

(123) No element, act, or instruction used herein should be construed as critical or essential unless explicitly described as such. Also, as used herein, the articles “a” and “an” are intended to include one or more items and may be used interchangeably with “one or more.” Further, as used herein, the article “the” is intended to include one or more items referenced in connection with the article “the” and may be used interchangeably with “the one or more.” Where only one item is intended, the phrase “only one,” “single,” or similar language is used. Also, as used herein, the terms “has,” “have,” “having,” or the like are intended to be open-ended terms that do not limit an element that they modify (e.g., an element “having” A may also have B). Further, the phrase “based on” is intended to mean “based, at least in part, on” unless explicitly stated otherwise. As used herein, the term “multiple” can be replaced with “a plurality of” and vice versa. Also, as used herein, the term “or” is intended to be inclusive when used in a series and may be used interchangeably with “and/or,” unless explicitly stated otherwise (e.g., if used in combination with “either” or “only one of”).

Claims

1. A memory device, comprising: a non-volatile cache memory; and one or more controllers configured to: receive, from a host device, an indication of a threshold amount of the non-volatile cache memory to be reserved for first data for which write boosting is activated; receive, from the host device, a write command that includes second data for which write boosting is deactivated;

determine that the second data is hot data; identify a block of the non-volatile cache memory using a write boost cursor based on determining that the second data is hot data, based on the indication of the threshold amount, and despite write boosting being deactivated for the second data; and write the second data to the block of the non-volatile cache memory.

2. The memory device of claim 1, wherein the one or more controllers, to identify the block using the write boost cursor, are configured to: identify the block using the write boost cursor based on a determination that writing the second data to the non-volatile cache memory will not reduce an available amount of non-volatile cache memory reserved for the first data below the threshold amount.

3. The memory device of claim 1, wherein the one or more controllers are further configured to copy hot data stored in the block to a first new block and copy cold data stored in the block to a second new block during a garbage collection operation performed on the block.

4. The memory device of claim 3, wherein the first new block is a first block of non-volatile main memory and the second new block is a second block of the non-volatile main memory.

5. The memory device of claim 3, wherein the first new block is a new block of the non-volatile cache memory and the second new block is a new block of non-volatile main memory.

6. The memory device of claim 3, wherein the one or more controllers are configured to copy the hot data to the first new block and copy the cold data to the second new block during the garbage collection operation based on one or more block age parameters associated with the non-volatile cache memory.

7. The memory device of claim 1, wherein the indication of the threshold amount comprises a threshold size or a threshold percentage.

8. The memory device of claim 1, further comprising storing small fragment data, having a size less than a page line size of a page line of the memory device, in the block of the non-volatile cache memory.

9. The memory device of claim 8, wherein a portion of the non-volatile cache memory is reserved for the small fragment data.

10. A method, comprising: receiving, by a memory device, a write command that includes data to be written to the memory device; receiving, by the memory device, an indication that single-level cell data caching is deactivated for the data; determining, by the memory device, whether the data is associated with a first data type or a second data type; and selectively writing, by the memory device, the data to single-level cell cache memory or multi-level cell main memory based on: a determination of whether the data is associated with the first data type or the second data type, and a determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching.

11. The method of claim 10, wherein the first data type is associated with an expected lifespan that is shorter than an expected lifespan associated with the second data type.

12. The method of claim 10, wherein selectively writing the data to the single-level cell cache memory or the multi-level cell main memory comprises writing the data to the single-level cell cache memory based on: a determination that the data is associated with the first data type, and a determination that the single-level cell cache memory has sufficient memory, that is not reserved for the single-level cell data caching, to store the data.

13. The method of claim 10, wherein selectively writing the data to the single-level cell cache memory or the multi-level cell main memory comprises writing the data to the multi-level cell main memory based on a determination that the data is associated with the second data type.

14. The method of claim 10, wherein selectively writing the data to the single-level cell cache memory or the multi-level cell main memory comprises writing the data to the multi-level cell main memory based on a determination that the single-level cell cache memory does not have sufficient memory, that is not reserved for the single-level cell data caching, to store the data.

15. The method of claim 10, further comprising receiving, from a host device, an indication of an

amount of memory in the single-level cell cache memory to be reserved for the single-level cell data caching, and wherein the determination of whether the single-level cell cache memory has available memory that is not reserved for the single-level cell data caching is based on the amount of memory in the single-level cell cache memory to be reserved for the single-level cell data caching.

16. The method of claim 10, further comprising: identifying a block of the single-level cell cache memory for which a garbage collection operation is to be performed; identifying a first set of pages, included in the block, that store first data associated with the first data type; identifying a second set of pages, included in the block, that store second data associated with the second data type; copying the first data to a first block as part of performing the garbage collection operation; and copying the second data to a second block as part of performing the garbage collection operation.

17. The method of claim 16, wherein both the first block and the second block are included in the multi-level cell main memory, or wherein the first block is included in the single-level cell cache memory and the second block is included in the multi-level cell main memory.

18. An apparatus, comprising: a non-volatile memory array; and a controller configured to execute instructions that cause the apparatus to: process a received command to write data to the non-volatile memory array, wherein the command is received without an indication that the data is to be written in a single-level cell (SLC) cache memory of the non-volatile memory array; determine whether the data is associated with a first data type; determine whether the SLC cache memory has sufficient memory not reserved for SLC data caching; and write the data to the SLC cache memory based on a determination that the data is associated with the first data type and that the SLC cache memory has sufficient memory not reserved for SLC data caching.

19. The apparatus of claim 18, wherein the controller is configured to execute instructions that cause the apparatus to: write the data to a multi-level cell memory of the non-volatile memory array based on a determination that the data is not associated with the first data type or that the SLC cache memory does not have sufficient memory not reserved for SLC data caching.

20. The apparatus of claim 18, wherein the first data type is expected to become invalid or is expected to be overwritten in a shorter amount of time than a second data type.
