

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386564
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Karkra; Kapil et al.

Reducing write amplification and over-provisioning using flash translation layer synchronization

Abstract

A host Flash Translation Layer (FTL) synchronizes host FTL operations with the drive FTL operations to reduce write amplification and over-provisioning. Embodiments of FTL synchronization map, at the host FTL software (SW) stack level, logical bands in which data is managed, referred to as host bands, to the physical bands on a drive where data is stored. The host FTL tracks validity levels of data managed in host bands to determine validity levels of data stored in corresponding physical bands, and optimizes defragmentation operations (such as garbage collection processes and trim operations) applied by the host FTL SW stack to the physical bands based on the tracked validity levels.

Inventors: Karkra; Kapil (Chandler, AZ), Malikowski; Wojciech (Gdansk, PL), Barczak; Mariusz (Gdansk, PL), Bahirat; Shirish (Longmont, CO)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 85176825

Appl. No.: 17/979687

Filed: November 02, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230051806 A1	Feb. 16, 2023

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

Field of Classification Search

CPC: G06F (3/0604); G06F (3/0607); G06F (3/0631); G06F (3/0665); G06F (3/0679); G06F (12/0246); G06F (2212/7201)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10592408	12/2019	Li et al.	N/A	N/A
10860228	12/2019	Mulani	N/A	G06F 3/0658
2013/0166826	12/2012	Franceschini	711/E12.008	G06F 12/0246
2018/0121344	12/2017	Seo	N/A	G06F 12/0246
2022/0404966	12/2021	Kanno	N/A	G06F 3/0619

OTHER PUBLICATIONS

Malikowski, Storage Performance Development Kit (SPDK) Flash Translation Layer Library for Zoned Namespace SSDs, Aug. 6, 2019, Flash Memory Summit (Year: 2019). cited by examiner Yang, Jingpei, et al., “Don't stack your Log on my Log”, SanDisk Corporation, Inflow'14, 2nd Workshop on Interactions of NVM/Flash with Operating Systems and Workloads, held Oct. 5, 2014, Broomfield, CO, USA., 10 pages. cited by applicant

Primary Examiner: Mackall; Larry T

Attorney, Agent or Firm: Haley Guiliano LLP

Background/Summary

TECHNICAL FIELD

(1) Embodiments described herein generally relate to an apparatus, computer program product, system, and method for managing data stored in memory devices, particularly NAND devices.

BACKGROUND

(2) Memory and storage devices are commonly used in computing systems for data storage and retrieval. To store and retrieve large amounts of data, servers and data centers in cloud computing or edge computing environments typically include solid state storage drives (SSDs) that use a non-volatile memory (NVM) Flash memory architecture (Flash) such as NAND (Negated “AND”) Flash memory components or devices.

(3) Because NAND Flash memory must be erased before new data can be written, SSDs that use NAND Flash memory as storage will have some data write amplification. Write amplification occurs when extra NAND operations are needed to move existing data, often more than once. These extra NAND operations produce a multiplying effect that increases the number of writes required, resulting in an “amplification” factor; hence the term “write amplification.” The write amplification factor (WAF) constantly changes over the life of the SSD.

(4) Write amplification has many implications for the read/write performance and the reliability of the SSD. Depending on the SSD's intelligence in managing its data, write amplification can cause

extra wear and extra read/write/erase cycles on the NAND components reducing the life of the NAND Flash component. An SSD's intelligence includes an on-drive Flash Translation Layer (FTL), typically implemented in the SSD firmware, which is designed to, among other things, reduce the amount of write amplification. In addition, an SSD is often "over-provisioned" with dedicated additional capacity to account for the extra read/write/erase cycles. Over-provisioned SSDs waste resources that could otherwise be used to increase storage capacity.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Features of embodiments of the claimed subject matter will become apparent as the following detailed description proceeds, and upon reference to the drawings, in which like numerals depict like parts, and in which:
- (2) FIG. 1 illustrates a block diagram of an example storage architecture in which an embodiment of FTL synchronization can be implemented.
- (3) FIG. 2 illustrates a block diagram of an example of write amplification data flow without the benefit of FTL synchronization.
- (4) FIG. 3A illustrates a block diagram of an example of the consequences of garbage collection performed without the benefit of FTL synchronization.
- (5) FIG. 3B illustrates a block diagram of an example of a write amplification that occurs without the benefit of FTL synchronization.
- (6) FIGS. 4-5 illustrate flow diagrams of example host FTL processes in accordance with embodiments FTL synchronization.
- (7) FIGS. 6A-6G illustrate flow diagrams of example host FTL write flow processes in accordance with embodiments FTL synchronization.
- (8) FIGS. 6H-6M illustrate flow diagrams of example host FTL garbage collection processes in accordance with embodiments FTL synchronization.
- (9) FIGS. 7-9 illustrate block diagrams of an example storage medium, storage device and computing platform that can be used to implement an embodiment of FTL synchronization.

DETAILED DESCRIPTION

- (10) To scale storage performance and storage capacity, some storage architectures include a host-based FTL (hereafter host FTL), typically implemented in a software (SW) stack, to manage write shaping, caching, data sequentialization and other operations designed to reduce write amplification. For example, a write shaping framework that sequences data using a host FTL SW stack on top of an on-drive FTL (hereafter referred to as a drive FTL) of a standard NAND SSD can artificially make a random workload look like sequential workload by absorbing write amplification at the host FTL SW stack.
- (11) Unfortunately, stacking a host FTL on top of one or more drive FTLs can have the unintended consequence of increasing write amplification instead of reducing write amplification for reasons that will be detailed in the description that follows.
- (12) To address this challenge, embodiments of FTL synchronization synchronize the host FTL operations with the drive FTL operations. Among other aspects, embodiments of FTL synchronization map, at the host FTL SW stack level, logical bands in which data is managed, referred to as host bands, to the physical bands on a drive where data is stored.
- (13) Among other information, embodiments of FTL synchronization are based in part on determining data validity levels. In NVM Flash devices, data is typically marked as valid or no longer valid by logical block address (LBA) as maintained in a logical-to-physical address table (L2P table). A host band validity level is typically expressed as a percentage of data managed in the host bands that is still valid. The validity level typically decreases over time as more of the data

gets erased or deleted from the physical band where it was stored, or is otherwise no longer valid. (14) In one embodiment FTL synchronization tracks host band validity levels of data managed in host bands. The host band validity levels of data managed in host bands allows the host FTL SW stack to indirectly determine validity levels of data stored in corresponding physical bands as will be described in further detail below. In this manner, FTL synchronization can optimize defragmentation operations (such as garbage collection processes and trim operations) applied by the host FTL SW stack to the physical bands based on the tracked host band validity levels.

(15) FIG. 1 illustrates a block diagram of an example storage architecture **100** in which an embodiment of FTL synchronization can be implemented. One or more applications **102** issue requests to create, write, modify and read data maintained in a storage system combining a high performance non-volatile memory (NVM) device **114** for caching and staging data, and one or more high capacity non-volatile memory (NVM) storage devices **118**, such as a NAND QLC (Quad Level Cell)/PLC (Penta Level Cell) SSD for storing data, hereafter referred to as an NVM storage drive **118**. The storage architecture **100** allows a host FTL **110** to receive user writes over an IO interface **104** and to access a write shaping buffer **108** and write shaping algorithms **112** to shape the user writes into shaped writes. The shaped writes are staged on the high performance NVM device **114** for writing back to the high capacity NVM storage drive **118**. The shaped writes are typically large, sequential, indirection-unit (IU) aligned writes designed to reduce write amplification at the drive FTL **120**.

(16) Despite the use of write shaping to reduce write amplification at the drive FTL **120**, the storage architecture **100** can actually cause an even greater write amplification problem by stacking the host FTL **110** on top of the drive FTL **120**. The reason is because the overall write amplification of the storage architecture **100** is multiplicative and not additive. For example, if the host FTL **110** write amplification factor (Host WAF) is 2 and drive FTL **120** write amplification factor (Drive WAF) is 3, then the overall system write amplification factor (System WAF) becomes $2 \times 3 = 6$.

(17) To examine the cause of this paradoxical increase in write amplification, FIG. 2 illustrates the write amplification data flow in detail **200**. As shown, user writes **212a/212b** are staged in a high performance NVM device **114**. In non-volatile cache **224**, write reduction process **202** and aggregation process **204** attempt to reduce writes to the high capacity NVM storage drive **118**. For example, the aggregation process **204** can combine small user writes **212a/212b** into large writes staged in the high performance NVM device **114**. A compaction process **206** issues compaction reads **214** to the high performance NVM device **114** and replaces invalid data (i.e., data marked as no longer valid) with valid data. The compaction process **206** further issues compaction writes **216** to a storage process **210** to store the compacted data to the high capacity NVM storage drive **118**. The host FTL **110** rewrites some of the data from the high capacity NVM storage drive **118** during the garbage collection process **208** using garbage collection reads **220** and garbage collection writes **218**. In addition, high capacity NVM storage drive **118** needs to eventually perform its own garbage collection process to make free space resulting in the drive FTL **120** issuing its own internal writes **222**. The overall write amplification factor (System WAF) from all of these writes can be computed using the following equation:

$$(\text{SystemWAF} = \text{HostWAF} * (1 - \text{WRF}) * (\text{DriveWAF}))$$

where System WAF represents the overall write amplification factor of storage architecture **100**,

(18) WRF represents the write reduction factor, Host WAF represents the write amplification attributable to the host FTL and Drive WAF represents the write amplification attributable to the drive FTL. The Host WAF can be computed using the following equation:

$$(19) \text{HostWAF} = \frac{\text{CompactionWrites} + \text{GarbageCollectionWrites}}{\text{CompactionWrites}}$$

The Drive WAF can be computed using the following equation:

$$(20) \text{DriveWAF} = 1 + \frac{\text{DriveInternalWrites}}{\text{UserWrites} * (1 - \text{WRF}) * \text{HostWAF}}$$

Lastly, the WRF can be computed using the following equation:

$$(21) \text{ WRF} = \frac{\text{UserWrites} - \text{CompactionWrites}}{\text{UserWrites}}$$

(22) Embodiments of FTL synchronization operate to minimize the overall System WAF by synchronizing the operation of the host FTL **110** with the drive FTL **120** through host band to physical band mapping, host band validity level tracking and host FTL **110** defragmentation operations performed on the high capacity NVM storage drive **118**, such as the garbage collection operations based on the tracked host band validity levels.

(23) In one embodiment, FTL synchronization operates to maintain the Drive WAF at or near unity, thereby eliminating or at least reducing the multiplicative effect on the overall System WAF. Thus, the above equation for System WAF

$$(\text{SystemWAF} = \text{HostWAF} * (1 - \text{WRF}) * (\text{DriveWAF}))$$

can be rewritten as:

$$(\text{SystemWAF} = \text{HostWAF} * (1 - \text{WRF}) * (1))$$

which simplifies to:

$$(\text{SystemWAF} = \text{HostWAF} * (1 - \text{WRF}))$$

(24) As a result, using embodiments of FTL synchronization as described herein, the overall System WAF is primarily driven by the Host WAF. For this reason, embodiments of FTL synchronization can eliminate, or at least reduce, the drive FTL over-provisioning to minimum acceptable levels in addition to reducing the Drive WAF.

(25) FIGS. 3A and 3B illustrate block diagrams of examples of the consequences of not keeping the Drive WAF in check using FTL synchronization. In FIG. 3A, example **300**, when host FTL **110** invokes garbage collection, it can try to find host bands (such as blocks or zones of data) which have the least valid data (minimal validity level) based on how the corresponding data is marked (valid or no longer valid) in a logical-to-physical address table (L2P) **116** maintained at the host FTL **110** SW stack level. However, the physical locations of the data in the high capacity NVM storage drive **118** might not be aligned to the media physical address granularity. For this reason, the host FTL **110** SW stack defragmentation operations may minimize the garbage collection for the host FTL **110** while unintentionally increasing the garbage collection for high capacity NVM storage drive **118**.

(26) This unintended consequence is illustrated by example **300**, in which optimum invalidated host bands, such as zones **302a**, **302b** and **302c** as determined by the host FTL **110** map to separate physical bands of the high capacity NVM storage drive **118**. An optimum invalidated host band contains data at a minimum validity level and is thus considered an optimal host band for garbage collection. Invalidated zone **302a** is on Physical Band **1** at location **304a**, invalidated zone **302b** is on Physical Band **2** at location **304b** and invalidated zone **302c** is on Physical Band **3** at location **304c**. Relocating the data that is still valid in those zones can increase the drive FTL **120** write amplification inside the high capacity NVM storage drive **118** in which the separate physical bands reside.

(27) Furthermore, in FIG. 3B, example **306**, when the host FTL **110** tries to optimize write amplification at the drive FTL **120** level directly by selecting a physical band on the high capacity NVM storage drive **118** with a minimum validity level for garbage collection, some of the data in the physical band can map to host bands (e.g., host-managed zones) that have high validity levels. As shown, the Physical Band **4** Invalid Zone **3** **310a**, Invalid Zone **2** **310b** and Invalid Zone **1** **310c**, as well as remaining Valid Zone **0** **312a**, Valid Zone a **312b**, Valid Zone b **312c**, and Valid Zone c **312d** contain valid data, all of which the host FTL **110** must transfer to the host-managed cache in Valid Zone x **308a**, Valid Zone y **308b**, Valid Zone z **308c**, and Valid zone k **308d**. Eventually, the host FTL **110** must re-write the transferred valid data back to the high capacity NVM storage drive **118**, thereby increasing the host FTL **110** write amplification.

(28) To address the foregoing challenges without the unintended consequences of increasing write amplification rather than decreasing it, embodiments of FTL synchronization can be implemented in the storage architecture **100** illustrated in FIG. **1** to provide physical band observability at the

host FTL **110** SW stack level for physical bands residing on the high capacity NVM storage drives **118**. Physical band observability allows the host FTL **110** to perform FTL operations that align with the physical bands. In addition, the host FTL **110** can optimize defragmentation operations by making any trim or unmap decisions at the host FTL **110** SW stack level using a single decision process instead of the current decoupled decision processes that occur across the host FTL **110** and the drive FTL **120**.

(29) Among other advantages, embodiments of FTL synchronization can eliminate or at least reduce drive FTL **120** write amplification because the host FTL **110** is able to invalidate an entire physical band inside the drive. This allows the elimination or reduction of over-provisioning inside the high capacity NVM storage drive **118** and a Drive WAF of 1, or at unity. Unlike other attempts to solve the problem of system write amplification, such as the NVMe Zoned Namespaces (ZNS) protocol in the NVMe 2.0 specification, embodiments of FTL synchronization can be implemented for all SSDs without having to adopt a different protocol.

(30) Depending on workload and requirements, the host FTL SW stack can tune its own over-provisioning and, with FTL synchronization, can reduce the host FTL SW stack over-provisioning to a minimum acceptable level. Importantly, since the host FTL **110** has the physical band observability of the drive FTL **120** internal physical band invalidation, the host FTL **110** can adapt the processes to reduce write amplification and over-provisioning accordingly. Lastly, as host bands and physical bands are aligned, the host FTL SW stack can optimize its block size or chunk size as a function of physical band capacity, thereby allowing FTL synchronization to be appropriately scaled to various capacities of a high capacity NVM storage drive **118**.

(31) Turning back to FIG. 1, in one embodiment, the host FTL **110** maintains a host band/physical band tracking data structure **106** that maps a host band identification (host band ID) **107a** representing a host band in which the host FTL **110** manages data with a physical band identification (physical band ID) **107b** representing a physical band on the one or more high capacity NVM storage drives **118** where data is stored. The host FTL **110** is capable of receiving a notification **122** from the one or more high capacity NVM storage drives **118** when a new physical band is opened. The notification **122** provides the host FTL **110** with the ability to map the host band IDs **107a** of the host bands in which data is managed to the physical band IDs **107b** of the physical bands where data is stored. In one embodiment, the notification **122** is transmitted as an Asynchronous Event Notification (AEN) in accordance with a communication protocol such as the Non-Volatile Memory Express (NVMe) Specification, revision 1.2, also published in November 2014 (“NVMe specification”).

(32) In one embodiment, the host band/physical band tracking data structure **106** further includes a host band status **107c** to indicate whether a host band is available to accommodate more data. Similar to status information for physical bands, the host band status **107c** indicates whether a newly opened host band contains no data (status=free), contains some data but available to accommodate more data (status=open), or is full and not available to accommodate more data (status=closed). The threshold for determining a host band status **107c** can be a threshold amount of host-managed space available to contain data, where the threshold can vary depending on the implementation.

(33) In one embodiment, the host band/physical band tracking data structure **106** further includes a host band validity level **107d**. As noted earlier, in NVM Flash devices, such as NVM storage device **118**, data is typically marked as valid or no longer valid by logical block address (LBA) as maintained in the L2P table **116**. A host band validity level **107d** is typically expressed as a percentage of data managed in the host band that is still valid. The validity level typically decreases over time as more of the data is marked as no longer valid due to being erased or deleted from the physical band where it was stored, or otherwise marked as no longer valid.

(34) FIGS. 4-5 illustrate flow diagrams of example host FTL write process **400** and an example host FTL garbage collection process **500** that can be used to implement an embodiment of FTL

synchronization as illustrated in FIG. 1. FIGS. 4-5 as illustrated herein provide examples of sequences of various process actions. The process diagrams can indicate processes, operations or actions to be executed by a software or firmware routine, as well as physical operations or actions. In one embodiment, a process diagram can illustrate the state of a finite state machine (FSM), which can be implemented in hardware and/or software. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes, operations or actions can be modified. Thus, the illustrated embodiments should be understood as an example, and the processes, operations or actions can be performed in a different order or in parallel. Additionally, one or more processes, operations or actions can be omitted in various embodiments; thus, not all processes, operations or actions are required in every embodiment. Other sequences or orders of processes, operations or actions are possible.

(35) In FIG. 4, the write process **400** begins at **402** in which a host FTL **110** selects a starting host band having a host band status of FREE indicating that the host band is empty, and changes the host band state to OPEN to reflect that the host band is available for managing data. At **404**, a host FTL **110** receives a user write from application **102**. At **406** the host FTL **110** writes the data to the selected host band. At **408**, the FTL **110** updates the host band validity level based on whether the data is marked as valid or no longer valid in the L2P table **116**. At **410**, the host FTL **110** maps the selected host band to the corresponding physical band where the data is stored.

(36) In one embodiment, at decision block **412**, the host FTL **110** determines whether the current selected host band is FULL, or if the corresponding physical band is FULL. If NO (neither is FULL), then processing continues to receive the next user write at **404**. If YES (either one or both are FULL), then at **414** the host FTL **110** changes the current selected host band state to CLOSED, and at **416**, concludes the process **400** by selecting the next FREE host band and changing the selected host band state to OPEN in preparation for receiving and processing the next user write at **404**.

(37) In FIG. 5, the garbage collection process begins at decision block **502** in which the host FTL **110** determines whether a threshold has been reached for the number of physical bands remaining in the FREE state. If NO, the host FTL **110** continues to monitor the number of physical bands remaining in the FREE state. But if YES, then the host FTL garbage collection process **500** continues at **504** to search for groups of one or more host bands that have been mapped to the same physical band where their data is stored. At **506** the process **500** continues in which the host FTL **110** selects from among the searched groups, the group that has the least validity based on an average validity level of the host bands in the group. The average validity level of the selected group can be used to determine that the corresponding physical band where the group's data is stored is a physical band with a minimum validity level, and thus a candidate for garbage collection.

(38) In one embodiment, at **508**, the host FTL **110** frees the physical band with the minimum validity by invalidating the entire physical band. The host FTL **110** moves the data that is still valid in the group's host bands to the next FREE or OPEN host bands that map to other physical bands that are also FREE or OPEN, such as the write process **400** in FIG. 4. The moves of valid data in the group's host bands is repeated until all valid data is moved. At **510**, the process **500** concludes with the host FTL **110** issuing trim commands to trim the remaining invalid content from the entire physical band, leaving it in a FREE state.

(39) FIGS. 6A-6G illustrate flow diagrams of example host FTL write flow processes in accordance with embodiments FTL synchronization. In FIGS. 6A-6G, a sequence of snapshots of an example write flow **600** for a host FTL **110** managing a high capacity storage drive **118** is shown. By way of example only, the high capacity storage drive **118** is a QLC 16 Terabyte drive with QLC physical band size of ~72GigaBytes (GB) each, and the host FTL **110** host band size is 1 GB. By way of example only, the host FTL **110** has a host band/physical band tracking data structure **602/106** designed in accordance with a Storage Performance Development Kit (SPDK)

Flash Translation Layer Library (<https://spdk.io/doc/ftl.html>). As shown in the data structure **602/106**, the physical band ID N starts from host band ID N*72 and ends at N*72+71. Based on this range formula, if N is 0 then the range is 0-71. If N is 1 then the range is 72 to 143, and so forth. If N is 253 then the range is 18216 to 18287. If N is 254, then the range is 18288 to 18359. And finally if N is 255 then the range is 18360 to 18431. The snapshots illustrate the host band states in the host band/physical band tracking data structure **602/106** and the physical band states **604** of the physical bands of the high capacity storage drive **118**. The types of physical band states and host band states are shown in the band state key **606**.

(40) The first snapshot of the example write flow **600** in FIG. **6A** shows that as the process commences, the host bands are free, and the first physical band ID **0** is open. Therefore, the FTL **110** selects host band ID **0** to write and change its state to OPEN. The second snapshot in FIG. **6B** shows that the first host band ID **0** is now full and corresponds to the first physical band ID **0**, which is still open. The third and fourth snapshots in FIG. **6C** and FIG. **6D** shows that the additional host band IDs through host band ID **71** have been filled and closed or opened, and the corresponding first physical band is still open but nearly full. As the host bands fill the host FTL changes the host band status to closed and selects the next available host band for writing. The fifth snapshot in FIG. **6E** shows that the host bands IDs **0** through **71** corresponding to physical band ID **0** are now all full and therefore the host FTL **110** has updated the host band status to closed. When the physical band is full, the host FTL **110** and/or drive FTL **120** changes the physical band status to closed as well. The sixth and seventh snapshots in FIGS. **6F** and **6G** show that the write flow **600** continues to open and fill the next available host bands and corresponding physical bands in order since the initial status of each host band and physical band is free. Although the host FTL **110** can select the next host band, it cannot control which physical band the host band data is written to. For ease of illustration and by way of example only, the next host band IDs **72** and **143** are shown as corresponding to physical band ID **2**, but the host band data could have been written to other open physical bands.

(41) FIGS. **6H-6M** illustrate flow diagrams of example host FTL garbage collection processes **608** in accordance with embodiments FTL synchronization. Continuing with the snapshots illustrated in FIGS. **6A-6G**, eventually the host FTL **110** will determine that only a few host bands are available for writing data and that the threshold for initiating a garbage collection process has been met. As shown in the eighth snapshot in FIG. **6H**, each host band ID has both a host status and a validity level expressed as a percentage of valid data remaining in the host bands identified by host band IDs. In one embodiment, the host FTL **110** searches for groups of host bands that are mapped to the same physical band in the host band/physical band tracking data structure **106**. As shown in FIG. **6H**, the group of host bands mapped to physical band ID **0** is determined to have an average validity level of 80% whereas the group corresponding to physical band ID **1** has an average validity level of 77%. Since 77% is the lower percentage value, the host FTL **110** commences moving the host bands corresponding to physical band **1** to the next available host band with a free or open status.

(42) The ninth and tenth snapshots in FIGS. **6I** and **6J** illustrate the effect of the host FTL **110** moving host band **72** to the next available host band 18288 corresponding to the beginning of physical band **254**, namely the freeing up of physical band **2**. The eleventh and twelfth snapshots in FIGS. **6K** and **6L** illustrate the continued freeing up of physical band **2**. After each move the host FTL **110** updates the host band statuses and validity levels to reflect the moves. Eventually, in the thirteenth snapshot in FIG. **6M**, the physical band **2** is now entirely free, with any remaining invalid data that was not moved to a new host band having been trimmed by the host FTL **110**. In this manner, the host FTL **110** is able to invalidate and free up an entire physical band on the high capacity NVM drive **118** without the involvement of the drive FTL **120**, thereby eliminating write amplification at the drive FTL **120** level during garbage collection and defragmentation operations.

(43) FIG. **7** illustrates an embodiment of a storage medium **700**. The storage medium **700** can

include an article of manufacture. In some examples, storage medium **700** can include any non-transitory computer readable medium or machine readable medium, such as an optical, magnetic or semiconductor storage. Storage medium **700** can store various types of computer readable and executable instructions **702**, such as instructions to implement the processes described herein. Examples of a computer readable or machine-readable storage medium can include any tangible media capable of storing electronic data, including volatile memory or non-volatile memory, removable or non-removable memory, erasable or non-erasable memory, writeable or re-writable memory, and so forth. Examples of computer readable and executable instructions **702** can include any suitable type of code, such as source code, compiled code, interpreted code, executable code, static code, dynamic code, object-oriented code, visual code, and the like.

(44) To the extent various operations or functions are described herein, they can be described or defined as software code, instructions, configuration, and/or data. The content can be directly executable (“object” or “executable” form), source code, or difference code (“delta” or “patch” code). The software content of the embodiments described herein can be provided via an article of manufacture with the content stored thereon, or via a method of operating a communication interface to send data via the communication interface. The machine readable storage medium can cause a machine to perform the functions or operations described, and includes any mechanism that stores information in a form accessible by a machine (e.g., computing device, electronic system, etc.), such as recordable/non-recordable media (e.g., read only memory (ROM), random access memory (RAM), magnetic disk storage media, optical storage media, flash memory devices, etc.).

(45) FIG. **8** illustrates an example storage device **800**. In some examples, as shown in FIG. **8**, storage device **800** can include a processing component **840**, storage device components **850**, and a communications interface **860**, as well as host a storage medium **700**. According to some examples, storage device **800** can be capable of being coupled to a host computing device or platform via the communications interface **860**.

(46) According to some examples, processing component **840** can execute processing operations or logic for a storage controller **810**, or for other components **850** of the storage device and storage medium **700**, and/or any activity described herein. Processing component **840** can include various hardware elements, firmware elements and software elements, or a combination of both.

(47) In some examples, storage device components **850** can include common computing elements or circuitry, such as one or more processors, multi-core processors, co-processors, memory units, chipsets, controllers, interfaces, oscillators, timing devices, power supplies, and so forth. Examples of memory units can include without limitation various types of computer readable and/or machine-readable storage media any other type of volatile or non-volatile storage media suitable for storing information.

(48) In some examples, communications interface **860** can include logic and/or features to support a communication interface. For these examples, communications interface **860** can include one or more communication interfaces that operate according to various communication protocols or standards to communicate over wired or wireless communication links. The communication interface **860** can include any mechanism that interfaces to any of a hardwired, wireless, optical, etc., medium to communicate to another device, such as a memory bus interface, a processor bus interface, an Internet connection, a disk controller, etc. The communication interface **860** can be configured by providing configuration parameters and/or sending signals to prepare the communication interface to provide a data signal describing the software content. The communication interface **860** can be accessed via one or more commands or signals sent to the communication interface.

(49) Storage device **800** can be configured as an SSD for use in storage architecture **100** as shown in FIG. **1**, including as a high performance NVM device **114** and/or a high capacity NVM device **118**. Accordingly, functions and/or specific configurations of storage device **800** described herein, can be included or omitted in various embodiments of storage device **800**, as suitably desired.

(50) FIG. 9 illustrates an example computing platform **900**. In some examples, as shown in FIG. 9, computing platform **900** can include one or more of: a storage system **930**, a processing component **940**, platform components **950**, or a communications interface **960**.

(51) According to some examples, storage system **930** can be used in the storage architecture **100** as shown in FIG. 1 and includes a controller **932** and memory device(s) **934**. For these examples, logic and/or features resident at or located at controller **932** can execute at least some processing operations or logic for the storage medium **700**, storage device **800**, or any other activities described herein. Also, memory device(s) **934** can include volatile or non-volatile memory devices that are described above for storage architecture **100**, including devices **114** and **118**.

(52) Non-volatile types of memory can be types of memory whose state is determinate even if power is interrupted to the device. In some examples, memory device(s) **934** can be block-addressable memory devices, such as memory devices including NAND or NOR technologies.

(53) According to some examples, volatile types of memory included in memory device(s) **934** can include, but are not limited to, random-access memory (RAM), Dynamic RAM (D-RAM), double data rate synchronous dynamic RAM (DDR SDRAM), static random-access memory (SRAM), thyristor RAM (T-RAM) or zero-capacitor RAM (Z-RAM). Volatile types of memory can be compatible with a number of memory technologies, such as DDR4 (DDR version 4, initial specification published in September 2012 by JEDEC), LPDDR4 (LOW POWER DOUBLE DATA RATE (LPDDR) version 4, JESD209-4, originally published by JEDEC in August 2014), WIO2 (Wide I/O2 (WideIO2), JESD229-2, originally published by JEDEC in August 2014), HBM (HIGH BANDWIDTH MEMORY DRAM, JESD235, originally published by JEDEC in October 2013), DDR5 (DDR version 5, currently in discussion by JEDEC), LPDDR5 (LPDDR version 5, currently in discussion by JEDEC), HBM2 (HBM version 2, currently in discussion by JEDEC), and/or others, and technologies based on derivatives or extensions of such specifications.

(54) According to some examples, processing component **940** can include various hardware elements, firmware elements, software elements, or a combination thereof. In some examples, platform components **950** can include common computing elements, such as one or more processors, single or multi-cores, co-processors, memory units, chipsets, controllers, peripherals, interfaces, oscillators, timing devices, video cards, audio cards, multimedia I/O components (e.g., digital displays), power supplies, and so forth. Examples of memory units associated with either platform components **950** or storage system **930** can include without limitation, various types of computer readable and machine-readable storage media.

(55) In some examples, communications interface **960** and communications interface **860**, can include logic and/or features to support a communication interface. For these examples, communications interface **860/960** can include one or more communication interfaces that operate according to various communication protocols or standards to communicate over direct or network communication links. Direct communications can occur through a direct interface via use of communication protocols or standards described in one or more industry standards (including progenies and variants) such as those associated with the SMBus specification, the PCIe specification, the NVMe specification, the SATA specification, SAS specification or the USB specification.

(56) Communications interface **860/960** can include radio frequency (RF) transceiver components for accessing wireless voice and/or data networks (e.g., using cellular telephone technology, data network technology such as 3G, 4G/LTE, Wi Fi, other IEEE 802.11 family standards, or other mobile communication technologies, or any combination thereof), components for short range wireless communication (e.g., using Bluetooth and/or Bluetooth LE standards, NFC, etc.), and/or other components. In some embodiments, communications interface **860/960** can provide wired network connectivity (e.g., Ethernet) in addition to or instead of a wireless interface.

Communications interface **860/960** can be implemented using a combination of hardware (e.g., driver circuits, antennas, modulators/demodulators, encoders/decoders, and other analog and/or

digital signal processing circuits) and software components. In some embodiments, communications interface **860/960** can support multiple communication channels concurrently or at different times, using the same transport or different transports.

(57) Computing platform **900** can be part of a computing device that can be, for example, user equipment, a computer, a personal computer (PC), a desktop computer, a laptop computer, a notebook computer, a netbook computer, a tablet, a smart phone, virtual reality or augment reality device, autonomous driving or flying vehicle, Internet-of-things (IoT) device, embedded electronics, a gaming console, a server, a server array or server farm, a web server, a network server, a proxy device, a work station, a mini-computer, a main frame computer, a supercomputer, a network appliance, a web appliance, a distributed computing system, multiprocessor systems, processor-based systems, or combination thereof. Accordingly, functions and/or specific configurations of computing platform **900** described herein, can be included or omitted in various embodiments of computing platform **900**, as suitably desired.

(58) The components and features of computing platform **900** can be implemented using any combination of discrete circuitry, ASICs, field programmable gate arrays (FPGAs), logic gates and/or single chip architectures. Further, the features of computing platform **900** can be implemented using microcontrollers, programmable logic arrays and/or microprocessors or any combination of the foregoing where suitably appropriate. It is noted that hardware, firmware and/or software elements can be collectively or individually referred to herein as “logic,” “circuit” or “circuitry.”

(59) Additional example implementations are as follows:

(60) Example 1 is a method, system, apparatus or computer-readable medium for flash translation layer (FTL) synchronization comprising a processor of a host device to execute a logic for a host FTL to manage data written to physical bands of a storage drive having a drive FTL, the logic including mapping a host band in which data is managed to a physical band where the data is stored, tracking a validity level for the host band representing how much of the data managed in the host band is valid, and optimizing defragmentation operations applied to the physical band by the host FTL based on tracking the validity level.

(61) Example 2 is the method, system, apparatus or computer-readable medium of Example 1 in which tracking the validity level is based on how much of the data managed in the host band is marked as valid in a logical-to-physical (L2P) table maintained in the host device.

(62) Example 3 is the method, system, apparatus or computer-readable medium of any of Examples 1 and 2 in which the logic further includes receiving, in the host device, a notification from the storage drive when physical bands are opened and mapping the host band in which the data is managed to an open physical band where the data is stored.

(63) Example 4 is the method, system, apparatus or computer-readable medium of any of Examples 1, 2 and 3 in which the logic further includes maintaining a host band state for the host band mapped to the physical band including updating the host band state to any of: FREE to indicate that the host band is empty and available to receive more data, OPEN to indicate that the host band is available to receive more data and CLOSED to indicate that the host band is any of full and not available to receive more data.

(64) Example 5 is the method, system, apparatus or computer-readable medium of any of Examples 1, 2, 3 and 4 in which the logic for tracking the validity level further includes: determining a group of host bands mapped to a same physical band, computing an average validity level based on an average of the validity level tracked for each host band in the group of host bands, and using the average validity level to represent a physical band validity level of the same physical band to which the group of host bands are mapped.

(65) Example 6 is the method, system, apparatus or computer-readable medium of any of Examples 1, 2, 3, 4 and 5, in which the logic for optimizing defragmentation operations applied to the physical band by the host FTL based on tracking the validity level further includes: determining the

average validity level is at or below a minimum validity level and removing all stored data from the same physical band to which the group of host bands are mapped, including: moving valid data to an available host band mapped to a different available physical band and trimming invalid data.

(66) Example 7 is the method, system, apparatus or computer-readable medium of any of Examples 1, 2, 3, 4, 5 and 6, in which the host band represents a logical space for managing data written to the physical band, the logical space including any one or more blocks, zones and chunks.

(67) Example 8 is a method, system, apparatus or computer-readable medium of any of Examples 1, 2, 3, 4, 5, 6 and 7 in which the logic for the host FTL is developed in accordance with a Storage Performance Development Kit (SPDK) FTL library.

(68) Example 9 is a method, system, apparatus or computer-readable medium of any of Examples 1, 2, 3, 4, 5, 6, 7 and 8, further comprising a controller to perform the logic for the host FTL.

(69) Example 10 is a method, system, apparatus or computer-readable medium for FTL synchronization comprising a non-volatile memory (NVM) storage drive and a processor of a host communicatively coupled to the NVM storage drive, in which the processor is to execute a host FTL to manage data in logical bands for storage in physical bands of the NVM storage drive, the host FTL to map a logical band of data in the host to a physical band in which the data is stored, track a validity level of the logical band indicating how much of the data managed in the logical band is valid, and perform garbage collection on the physical band based on a tracked validity level of one or more logical bands mapped to the physical band.

(70) Example 11 is the method, system, apparatus or computer-readable medium of Example 10 in which the host FTL is further to: receive a notification from a drive FTL of the NVM storage drive when a physical band is opened and map a logical band of data in the host to an opened physical band in which the data is stored.

(71) Example 12 is the method, system, apparatus or computer-readable medium of any of Example 10 and Example 11, in which the logical bands include any one or more of blocks, zones and chunks of a host-managed space and the host FTL is developed in accordance with a Storage Performance Development Kit (SPDK) FTL library.

(72) Example 13 is the method, system, apparatus or computer-readable medium of any of Examples 10, 11 and 12 in which to track the validity level of the logical band the processor is further to: determine a group of logical bands mapped to a same physical band, compute an average validity level based on an average of the tracked validity level for each logical band in the group and use the average validity level to represent a physical band validity level of the same physical band to which the group is mapped.

(73) Example 14 is the method, system, apparatus or computer-readable medium of any of Examples 10, 11, 12 and 13, in which to perform garbage collection on the physical band based on the tracked validity level the processor is further to: determine the average validity level is at or below a minimum validity level and remove all stored data from the same physical band to which the group of logical bands is mapped, including to move valid data to an available logical band mapped to a different available physical band and trim invalid data.

(74) The appearances of the phrase “one example” or “an example” are not necessarily all referring to the same example or embodiment. Any aspect described herein can be combined with any other aspect or similar aspect described herein, regardless of whether the aspects are described with respect to the same figure or element.

(75) Some examples can be described using the expression “coupled” and “connected” along with their derivatives. These terms are not necessarily intended as synonyms for each other. For example, descriptions using the terms “connected” and/or “coupled” can indicate that two or more elements are in direct physical or electrical contact with each other. The term “coupled,” however, can also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

(76) The terms “first,” “second,” and the like, herein do not denote any order, quantity, or

importance, but rather are used to distinguish one element from another. The terms “a” and “an” herein do not denote a limitation of quantity, but rather denote the presence of at least one of the referenced items. The term “asserted” used herein with reference to a signal denote a state of the signal, in which the signal is active, and which can be achieved by applying any logic level either logic 0 or logic 1 to the signal. The terms “follow” or “after” can refer to immediately following or following after some other event or events.

(77) Disjunctive language such as the phrase “at least one of X, Y, or Z,” unless specifically stated otherwise, is otherwise understood within the context as used in general to present that an item, term, etc., can be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present. Additionally, conjunctive language such as the phrase “at least one of X, Y, and Z,” unless specifically stated otherwise, should also be understood to mean X, Y, Z, or any combination thereof, including “X, Y, and/or Z.”

(78) Besides what is described herein, various modifications can be made to the disclosed embodiments and implementations of the invention without departing from their scope.

(79) Therefore, the illustrations and examples herein should be construed in an illustrative, and not a restrictive sense. The scope of the invention should be measured solely by reference to the claims that follow.

Claims

1. A computer-implemented method of comprising: in a processor of a host device to execute a logic for a host flash translation layer (FTL) to manage data written to physical bands of a storage drive having a drive FTL, the logic including: mapping a host band in which data is managed to a physical band where the data is stored; tracking a validity level for the host band representing how much of the data managed in the host band is valid; optimizing defragmentation operations applied to the physical band by the host FTL based on tracking the validity level; determining a group of host bands mapped to a same physical band; computing an average validity level based on an average of the validity level tracked for each host band in the group of host bands; and using the average validity level to represent a physical band validity level of the same physical band to which the group of host bands are mapped.
2. The computer-implemented method of claim 1, wherein tracking the validity level is based on how much of the data managed in the host band is marked as valid in a logical-to-physical (L2P) table maintained in the host device.
3. The computer-implemented method of claim 1, the logic further including: receiving, in the host device, a notification from the storage drive when physical bands are opened; and mapping the host band in which the data is managed to an open physical band where the data is stored.
4. The computer-implemented method of claim 1, the logic further including: maintaining a host band state for the host band mapped to the physical band including updating the host band state to any of: FREE to indicate that the host band is empty and available to receive more data; OPEN to indicate that the host band is available to receive more data; and CLOSED to indicate that the host band is any of full and not available to receive more data.
5. The computer-implemented method of claim 1, wherein the logic for optimizing defragmentation operations applied to the physical band by the host FTL based on tracking the validity level further includes: determining the average validity level is at or below a minimum validity level; and removing all stored data from the same physical band to which the group of host bands are mapped, including: moving valid data to an available host band mapped to a different available physical band, and trimming invalid data.
6. The computer-implemented method of claim 1, wherein the host band represents a logical space

for managing data written to the physical band, the logical space including any one or more blocks, zones and chunks.

7. The computer-implemented method of claim 1, wherein the logic for the host FTL is developed in accordance with a Storage Performance Development Kit (SPDK) FTL library.

8. An apparatus comprising: a host having a host flash translation layer (FTL) to manage how data is written to a storage drive having a drive FTL; a controller to perform the host FTL, including to: map a host band in which data is managed to a physical band of the storage drive where the data is stored, track a validity level for the host band representing how much of the data managed in the host band is valid, and optimize defragmentation operations applied to the physical band by the host FTL based on a tracked validity level for the host band; and wherein to track the validity level, the controller further to: determine a group of host bands mapped to a same physical band, compute an average validity level based on an average of the tracked validity level for each host band in the group of host bands, and use the average validity level to represent a physical band validity level of the same physical band to which the group of host bands are mapped.

9. The apparatus as in claim 8, wherein the tracked validity level is based on how much of the data managed in the host band is marked as valid in a logical-to-physical (L2P) table maintained in the host.

10. The apparatus as in claim 8, wherein to perform the host FTL the controller further to: receive, at the host, a notification from the storage drive when physical bands are opened; map the host band in which data is managed to an open physical band where the data is stored; and maintain a host band state for the host band mapped to the physical band including to update the host band state to any of: FREE to indicate that the host band is empty and available to receive more data, OPEN to indicate that the host band is available to receive more data, and CLOSED to indicate that the host band is any of full and not available to receive more data.

11. The apparatus as in claim 8, wherein the host band represents a logical space for managing data written to the physical band, the logical space including any one or more blocks, zones and chunks.

12. The apparatus as in claim 8, wherein the host FTL is developed in accordance with a Storage Performance Development Kit (SPDK) FTL library.

13. The apparatus as in claim 8, wherein to optimize defragmentation operations applied to the physical band by the host FTL the controller further to: determine the average validity level is at or below a minimum validity level; and remove all stored data from the same physical band to which the group of host bands are mapped, including: to move valid data to an available host band mapped to a different available physical band, and to trim invalid data.

14. A system comprising: a non-volatile memory (NVM) storage drive; a processor of a host communicatively coupled to the NVM storage drive, the processor to execute a host flash translation layer (FTL) to manage data in logical bands for storage in physical bands of the NVM storage drive, the host FTL to: map a logical band of data in the host to a physical band in which the data is stored, track a validity level of the logical band indicating how much of the data managed in the logical band is valid, and perform garbage collection on the physical band based on a tracked validity level of one or more logical bands mapped to the physical band; and wherein to track the validity level of the logical band the processor is further to: determine a group of logical bands mapped to a same physical band, compute an average validity level based on an average of the tracked validity level for each logical band in the group, and use the average validity level to represent a physical band validity level of the same physical band to which the group is mapped.

15. The system as in claim 14, the host FTL further to: receive a notification from a drive FTL of the NVM storage drive when a physical band is opened; and map a logical band of data in the host to an opened physical band in which the data is stored.

16. The system as in claim 14, wherein: the logical bands include any one or more of blocks, zones and chunks of a host-managed space; and the host FTL is developed in accordance with a Storage Performance Development Kit (SPDK) FTL library.

17. The system as in claim 14, wherein to perform garbage collection on the physical band based on the tracked validity level, the processor is further to: determine the average validity level is at or below a minimum validity level; and remove all stored data from the same physical band to which the group of logical bands is mapped, including: moving valid data to an available logical band mapped to a different available physical band, and trimming invalid data.
