

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386585
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Ballinger; Brandon M. et al.

Multi-modal input on an electronic device

Abstract

A computer-implemented input-method editor process includes receiving a request from a user for an application-independent input method editor having written and spoken input capabilities, identifying that the user is about to provide spoken input to the application-independent input method editor, and receiving a spoken input from the user. The spoken input corresponds to input to an application and is converted to text that represents the spoken input. The text is provided as input to the application.

Inventors: Ballinger; Brandon M. (San Francisco, CA), Schalkwyk; Johan (Scarsdale, NY), Cohen; Michael H. (Portola Valley, CA), Byrne; William J. (Davis, CA), Hafsteinsson; Gudmundur (Los Gatos, CA), Lebeau; Michael J. (New York, NY)

Applicant: Google LLC (Mountain View, CA)

Family ID: 43858442

Assignee: Google LLC (Mountain View, CA)

Appl. No.: 18/421189

Filed: January 24, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240160403 A1	May. 16, 2024

Related U.S. Application Data

continuation parent-doc US 17812320 20220713 US 11914925 child-doc US 18421189
continuation parent-doc US 16892749 20200604 US 11416214 20220816 child-doc US 17812320
continuation parent-doc US 16169279 20181024 US 10713010 20200714 child-doc US 16892749
continuation parent-doc US 14988408 20160105 US 10157040 20181218 child-doc US 16169279
continuation parent-doc US 14299837 20140609 US 9251791 20160202 child-doc US 14988408

continuation parent-doc US 13249172 20110929 US 8751217 20140610 child-doc US 14299837
continuation parent-doc US 12977003 20101222 US 9031830 20150512 child-doc US 13249172
us-provisional-application US 61330219 20100430
us-provisional-application US 61289968 20091223

Publication Classification

Int. Cl.: **G10L15/183** (20130101); **G06F3/04886** (20220101); **G06F3/16** (20060101);
G06F40/284 (20200101); **G06F40/58** (20200101); **G10L15/00** (20130101); **G10L15/18**
(20130101); **G10L15/22** (20060101); **G10L15/26** (20060101); **G10L15/30** (20130101);
G10L15/197 (20130101)

U.S. Cl.:

CPC **G06F3/167** (20130101); **G06F3/04886** (20130101); **G06F40/284** (20200101);
G06F40/58 (20200101); **G10L15/005** (20130101); **G10L15/18** (20130101);
G10L15/183 (20130101); **G10L15/22** (20130101); **G10L15/26** (20130101); **G10L15/30**
(20130101); G10L15/197 (20130101); G10L2015/223 (20130101); G10L2015/228
(20130101)

Field of Classification Search

CPC: G10L (15/22); G10L (15/005); G10L (15/18); G10L (15/183); G10L (15/26); G10L
(15/30); G10L (2015/223); G10L (2015/228); G10L (15/197)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
4820059	12/1988	Miller et al.	N/A	N/A
5267345	12/1992	Brown et al.	N/A	N/A
5632002	12/1996	Hashimoto et al.	N/A	N/A
5638487	12/1996	Chigier	N/A	N/A
5715367	12/1997	Gillick et al.	N/A	N/A
5737724	12/1997	Atal et al.	N/A	N/A
5768603	12/1997	Brown et al.	N/A	N/A
5805832	12/1997	Brown et al.	N/A	N/A
5822730	12/1997	Roth et al.	N/A	N/A
6021403	12/1999	Horvitz et al.	N/A	N/A
6119186	12/1999	Watts et al.	N/A	N/A
6167377	12/1999	Gillick et al.	N/A	N/A
6182038	12/2000	Balakrishnan et al.	N/A	N/A
6317712	12/2000	Kao et al.	N/A	N/A
6356866	12/2001	Pratley et al.	N/A	N/A
6397180	12/2001	Jaramillo et al.	N/A	N/A
6418431	12/2001	Mahajan et al.	N/A	N/A
6430551	12/2001	Thelen et al.	N/A	N/A
6446041	12/2001	Reynar et al.	N/A	N/A
6539358	12/2002	Coon et al.	N/A	N/A

6581033	12/2002	Reynar et al.	N/A	N/A
6678415	12/2003	Popat et al.	N/A	N/A
6714778	12/2003	Nykanen et al.	N/A	N/A
6778959	12/2003	Wu et al.	N/A	N/A
6822585	12/2003	Ni et al.	N/A	N/A
6839670	12/2004	Stammmler et al.	N/A	N/A
6851115	12/2004	Cheyer et al.	N/A	N/A
6876966	12/2004	Deng et al.	N/A	N/A
6912499	12/2004	Sabourin et al.	N/A	N/A
6922669	12/2004	Schalk et al.	N/A	N/A
6950796	12/2004	Ma et al.	N/A	N/A
6959276	12/2004	Droppo et al.	N/A	N/A
6985147	12/2005	Asakawa et al.	N/A	N/A
7027987	12/2005	Franz et al.	N/A	N/A
7043422	12/2005	Gao et al.	N/A	N/A
7107204	12/2005	Liu et al.	N/A	N/A
7149688	12/2005	Schalkwyk	N/A	N/A
7149970	12/2005	Pratley et al.	N/A	N/A
7174288	12/2006	Ju et al.	N/A	N/A
7200550	12/2006	Menezes et al.	N/A	N/A
7203731	12/2006	Coates et al.	N/A	N/A
7257532	12/2006	Toyama	N/A	N/A
7310601	12/2006	Nishizaki et al.	N/A	N/A
7328155	12/2007	Endo et al.	N/A	N/A
7366668	12/2007	Franz et al.	N/A	N/A
7370275	12/2007	Haluptzok et al.	N/A	N/A
7383553	12/2007	Atkin et al.	N/A	N/A
7392188	12/2007	Junkawitsch et al.	N/A	N/A
7403888	12/2007	Wang et al.	N/A	N/A
7424426	12/2007	Furui et al.	N/A	N/A
7424428	12/2007	Rose et al.	N/A	N/A
7451085	12/2007	Rose et al.	N/A	N/A
7493260	12/2008	Harb et al.	N/A	N/A
7505894	12/2008	Menezes et al.	N/A	N/A
7526431	12/2008	Roth et al.	N/A	N/A
7577562	12/2008	Menezes et al.	N/A	N/A
7624018	12/2008	Chambers et al.	N/A	N/A
7634720	12/2008	Haluptzok et al.	N/A	N/A
7672833	12/2009	Blume et al.	N/A	N/A
7698124	12/2009	Menezes et al.	N/A	N/A
7698136	12/2009	Nguyen et al.	N/A	N/A
7752046	12/2009	Bacchiani et al.	N/A	N/A
7778816	12/2009	Reynar	N/A	N/A
7805299	12/2009	Coifman	N/A	N/A
7818279	12/2009	Liu et al.	N/A	N/A
7831427	12/2009	Potter et al.	N/A	N/A
7848927	12/2009	Ohno et al.	N/A	N/A
7849080	12/2009	Chang et al.	N/A	N/A
7856351	12/2009	Yaman et al.	N/A	N/A
7877385	12/2010	Craswell et al.	N/A	N/A

7881936	12/2010	Longe et al.	N/A	N/A
7890326	12/2010	Strope et al.	N/A	N/A
7907705	12/2010	Huff et al.	N/A	N/A
7921107	12/2010	Chang et al.	N/A	N/A
7941189	12/2010	Miyauchi	N/A	N/A
7953692	12/2010	Bower et al.	N/A	N/A
7996220	12/2010	Rose et al.	N/A	N/A
8001130	12/2010	Wen et al.	N/A	N/A
8005680	12/2010	Kommer	N/A	N/A
8009678	12/2010	Brooke	N/A	N/A
8027973	12/2010	Cao et al.	N/A	N/A
8060373	12/2010	Gibbon et al.	N/A	N/A
8069027	12/2010	Liu et al.	N/A	N/A
8069043	12/2010	Bacchiani et al.	N/A	N/A
8073692	12/2010	Ativanichayaphong et al.	N/A	N/A
8156129	12/2011	Zhou et al.	N/A	N/A
8195319	12/2011	Kristjansson	455/414.4	G06F 3/167
8224656	12/2011	Scholz et al.	N/A	N/A
8620667	12/2012	Andrew	N/A	N/A
8670987	12/2013	Bergl et al.	N/A	N/A
8751217	12/2013	Ballinger	704/251	G10L 15/18
8775154	12/2013	Clinchant et al.	N/A	N/A
8813060	12/2013	Tewari et al.	N/A	N/A
8847775	12/2013	Morin et al.	N/A	N/A
8949266	12/2014	Phillips et al.	N/A	N/A
9031830	12/2014	Ballinger	704/251	G06F 3/04886
9165257	12/2014	Badger	N/A	G06F 16/35
9251791	12/2015	Ballinger	N/A	G06F 3/167
9495127	12/2015	Ballinger	N/A	G06F 40/284
9535906	12/2016	Lee	N/A	G06F 3/0489
10157040	12/2017	Ballinger	N/A	G10L 15/26
10713010	12/2019	Ballinger	N/A	G10L 15/22
11416214	12/2021	Ballinger	N/A	G06F 40/58
11914925	12/2023	Ballinger	N/A	G10L 15/26
2002/0055844	12/2001	L'Esperance	704/E15.045	G06F 3/165
2002/0062216	12/2001	Guenther et al.	N/A	N/A
2002/0087309	12/2001	Lee et al.	N/A	N/A

2002/0087314	12/2001	Fischer et al.	N/A	N/A
2002/0087325	12/2001	Lee et al.	N/A	N/A
2002/0099818	12/2001	Russell et al.	N/A	N/A
2002/0111990	12/2001	Wood et al.	N/A	N/A
2003/0050778	12/2002	Nguyen et al.	N/A	N/A
2003/0149561	12/2002	Zhou	N/A	N/A
2003/0216919	12/2002	Roushar	N/A	N/A
2003/0236099	12/2002	Deisher et al.	N/A	N/A
2004/0024583	12/2003	Freeman	N/A	N/A
2004/0043758	12/2003	Sorvari et al.	N/A	N/A
2004/0049388	12/2003	Roth et al.	N/A	N/A
2004/0098571	12/2003	Falcon	N/A	N/A
2004/0138882	12/2003	Miyazawa	N/A	N/A
2004/0172258	12/2003	Dominach et al.	N/A	N/A
2004/0230420	12/2003	Kadambe et al.	N/A	N/A
2004/0243415	12/2003	Commarford et al.	N/A	N/A
2005/0005240	12/2004	Reynar et al.	N/A	N/A
2005/0060138	12/2004	Wang et al.	N/A	N/A
2005/0091037	12/2004	Haluptzok et al.	N/A	N/A
2005/0108017	12/2004	Esser et al.	N/A	N/A
2005/0114474	12/2004	Anderson et al.	N/A	N/A
2005/0187763	12/2004	Arun	N/A	N/A
2005/0193144	12/2004	Hassan et al.	N/A	N/A
2005/0203740	12/2004	Chambers et al.	N/A	N/A
2005/0216273	12/2004	Reding et al.	N/A	N/A
2005/0246325	12/2004	Pettinati et al.	N/A	N/A
2005/0283364	12/2004	Longe et al.	N/A	N/A
2006/0004572	12/2005	Ju et al.	N/A	N/A
2006/0004850	12/2005	Chowdhury	N/A	N/A
2006/0009974	12/2005	Junqua et al.	N/A	N/A
2006/0035632	12/2005	Sorvari et al.	N/A	N/A
2006/0048055	12/2005	Wu et al.	N/A	N/A
2006/0095248	12/2005	Menezes et al.	N/A	N/A
2006/0100876	12/2005	Nishizaki et al.	N/A	N/A
2006/0111891	12/2005	Menezes et al.	N/A	N/A
2006/0111892	12/2005	Menezes et al.	N/A	N/A
2006/0111896	12/2005	Menezes et al.	N/A	N/A
2006/0111907	12/2005	Mowatt et al.	N/A	N/A
2006/0148520	12/2005	Baker et al.	N/A	N/A
2006/0212288	12/2005	Sethy et al.	N/A	N/A
2006/0217990	12/2005	Theimer et al.	N/A	N/A
2006/0247915	12/2005	Bradford et al.	N/A	N/A
2007/0060114	12/2006	Ramer et al.	N/A	N/A
2007/0078822	12/2006	Cucerzan et al.	N/A	N/A
2007/0106685	12/2006	Houh et al.	N/A	N/A
2007/0124507	12/2006	Gurram et al.	N/A	N/A
2007/0156407	12/2006	Schedl	N/A	N/A
2007/0174040	12/2006	Liu et al.	N/A	N/A
2007/0214131	12/2006	Cucerzan et al.	N/A	N/A
2007/0233488	12/2006	Carus et al.	N/A	N/A

2007/0260972	12/2006	Anderl	N/A	N/A
2008/0005067	12/2007	Dumais et al.	N/A	N/A
2008/0027723	12/2007	Reding et al.	N/A	N/A
2008/0030976	12/2007	Murazaki et al.	N/A	N/A
2008/0091406	12/2007	Baldwin et al.	N/A	N/A
2008/0091435	12/2007	Strope et al.	N/A	N/A
2008/0091443	12/2007	Strope et al.	N/A	N/A
2008/0131851	12/2007	Kanevsky et al.	N/A	N/A
2008/0133228	12/2007	Rao	N/A	N/A
2008/0155111	12/2007	Takeuchi et al.	N/A	N/A
2008/0162138	12/2007	Gurram et al.	N/A	N/A
2008/0174570	12/2007	Jobs	345/173	H04M 1/72403
2008/0188271	12/2007	Miyauchi	N/A	N/A
2008/0215541	12/2007	Li et al.	N/A	N/A
2008/0221887	12/2007	Rose et al.	N/A	N/A
2008/0221898	12/2007	Cerra et al.	N/A	N/A
2008/0221902	12/2007	Cerra et al.	N/A	N/A
2008/0228903	12/2007	Davies et al.	N/A	N/A
2008/0301112	12/2007	Wu	N/A	N/A
2008/0312910	12/2007	Zhang	N/A	N/A
2009/0018829	12/2008	Kuperstein	N/A	N/A
2009/0018833	12/2008	Kozat et al.	N/A	N/A
2009/0030687	12/2008	Cerra et al.	N/A	N/A
2009/0030696	12/2008	Cerra	704/275	G10L 15/30
2009/0055162	12/2008	Qian et al.	N/A	N/A
2009/0063455	12/2008	Li et al.	N/A	N/A
2009/0150160	12/2008	Mozer	N/A	N/A
2009/0164216	12/2008	Chengalvarayan et al.	N/A	N/A
2009/0170561	12/2008	Chou	N/A	N/A
2009/0210214	12/2008	Qian et al.	N/A	N/A
2009/0216531	12/2008	Yanagihara	704/235	G10L 15/30
2009/0254346	12/2008	Moore et al.	N/A	N/A
2009/0271177	12/2008	Menezes et al.	N/A	N/A
2009/0271188	12/2008	Agapi et al.	N/A	N/A
2009/0276219	12/2008	Yeh et al.	N/A	N/A
2009/0287681	12/2008	Paek et al.	N/A	N/A
2009/0287842	12/2008	Plamondon	N/A	N/A
2009/0292529	12/2008	Bangalore et al.	N/A	N/A
2009/0299745	12/2008	Kennewick et al.	N/A	N/A
2010/0004930	12/2009	Strope et al.	N/A	N/A
2010/0042400	12/2009	Block et al.	N/A	N/A
2010/0049502	12/2009	Oppenheim et al.	N/A	N/A
2010/0057687	12/2009	Shen et al.	N/A	N/A
2010/0088303	12/2009	Chen et al.	N/A	N/A
2010/0100377	12/2009	Madhavapeddi	704/235	G10L 15/22
2010/0106497	12/2009	Phillips	N/A	N/A

2010/0111059	12/2009	Bappu et al.	N/A	N/A
2010/0153219	12/2009	Mei et al.	N/A	N/A
2010/0153370	12/2009	Gollapudi et al.	N/A	N/A
2010/0179803	12/2009	Sawaf et al.	N/A	N/A
2010/0211588	12/2009	Jiang et al.	N/A	N/A
2010/0217581	12/2009	Hong	N/A	N/A
2010/0254521	12/2009	Kriese et al.	N/A	N/A
2010/0257171	12/2009	Shekhawat	N/A	N/A
2010/0257478	12/2009	Longe et al.	N/A	N/A
2010/0261526	12/2009	Anderson	345/157	G06F 3/016
2010/0286979	12/2009	Zangvil et al.	N/A	N/A
2010/0315266	12/2009	Gunawardana et al.	N/A	N/A
2010/0318531	12/2009	Gao et al.	N/A	N/A
2010/0325109	12/2009	Bai et al.	N/A	N/A
2011/0004462	12/2010	Houghton et al.	N/A	N/A
2011/0060587	12/2010	Phillips	704/235	G10L 15/30
2011/0066577	12/2010	Van Gael et al.	N/A	N/A
2011/0077943	12/2010	Miki et al.	N/A	N/A
2011/0093265	12/2010	Stent et al.	N/A	N/A
2011/0137653	12/2010	Ljolje et al.	N/A	N/A
2011/0161072	12/2010	Terao et al.	N/A	N/A
2011/0161081	12/2010	Ballinger et al.	N/A	N/A
2011/0162035	12/2010	King et al.	N/A	N/A
2011/0231183	12/2010	Yamamoto et al.	N/A	N/A
2011/0238410	12/2010	Larcheveque et al.	N/A	N/A
2011/0257974	12/2010	Kristjansson et al.	N/A	N/A
2011/0288868	12/2010	Lloyd et al.	N/A	N/A
2011/0295590	12/2010	Lloyd et al.	N/A	N/A
2011/0295852	12/2010	Wang et al.	N/A	N/A
2012/0022867	12/2011	Ballinger et al.	N/A	N/A
2012/0022868	12/2011	LeBeau et al.	N/A	N/A
2012/0060113	12/2011	Sejnoha et al.	N/A	N/A
2012/0143611	12/2011	Qian et al.	N/A	N/A
2012/0191745	12/2011	Velipasaoglu et al.	N/A	N/A
2012/0296655	12/2011	Kristjansson	704/E15.001	G06F 1/1694
2013/0046544	12/2012	Kay	345/169	G06F 3/023
2014/0173440	12/2013	Dal Mutto et al.	N/A	N/A
2014/0324434	12/2013	Vozila et al.	N/A	N/A
2024/0160403	12/2023	Ballinger	N/A	G10L 15/22

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
10045020	12/2000	DE	N/A
2101250	12/2008	EP	N/A

2002096070	12/2001	WO	N/A
2005122143	12/2004	WO	N/A

OTHER PUBLICATIONS

Extended European Search Report in European Application No. 16001249.8-1901/3091535 dated Jan. 31, 2017, 9 pages. cited by applicant

Bocchieri et al. "Use of geographical meta-data in ASR language and acoustic models," Acoustics Speech and Signal Processing (ICASSP), 2010 IEEE International Conference on IEEE, Mar. 14, 2010, pp. 5118-5121. cited by applicant

Cao et al. "Context-aware query suggestion by mining click-through and session data," in Proceedings of the 14th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '08), ACM, 2008, 875-883. cited by applicant

Dhillon. "Co-clustering documents and words using bipartite spectral graph partitioning:" in Proceedings of the Seventh ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, (KDD '01), ACM, New York, NY, 2001, 269-274. cited by applicant

Frey et al. "ALGONQUIN: Iterating Laplace's Method to Remove Multiple Types of Acoustic Distortion for Robust Speech Recognition", Eurospeech 2001 Scandinavia, 7th European Conference on Speech Communication and Technology, Aalborg, Denmark, Sep. 3-7, 2001, [online]. research.microsoft.com [retrieved on Nov. 23, 2010]. Retrieved from the Internet: <URL: <<http://research.microsoft.com/pubs/76516/2001-frev-eurospeech.pdf>> 4 pages. cited by applicant

"Japanese Input on Windows XP," available at http://blog.greggman.com/blog/japanese_input_on_windows_xp, Apr. 29, 2004, 11 pages. cited by applicant

Kristjansson et al. "Super-Human Multi-Talker Speech Recognition: The IBM 2006 Speech Separation Challenge System", Interspeech 2006: ICSLP; Proceedings of the Ninth International Conference on Spoken Language Process, Sep. 17-21, 2006, Pittsburgh, Pennsylvania, USA, Bonn, ISCA, 2006, 4 pages. cited by applicant

Kurland. Re-ranking search results using language models of query-specific clusters, ACM, 2009, 12(4):437-460. cited by applicant

Lee et al. "Search Result Clustering Using Label Language Model," IJCNLP 2008, The Third International Joint Conference on Natural Language Processing, Jan. 7-12, 2008, Hyderabad, India, 6 pages. cited by applicant

Li et al. "Query-URL Bipartite Based Approach to Personalized Query Recommendation," Proceedings of the Twenty Third AAAI Conference on Artificial Intelligence, 2008, 6 pages. cited by applicant

Liu and Croft. "Cluster-based retrieval using language models:" in Proceedings of the 27th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, (SIGIR '04). ACM, New York, NY, 2004, 186-193. cited by applicant

Maguitman et al. "Algorithmic detection of semantic similarity," in Proceedings of the 14th International Conference on World Wide Web (WWW '05), ACM, 2005, 107-116. cited by applicant

Mohri et al. "Weighted Finite-State Transducers in Speech Recognition," Computer Speech & Language, 2002, 16 (1):69-88. cited by applicant

Vertanen. "An Overview of Discriminative Training for Speech Recognition," Technical Report, 2004, from http://www.inference.phy.cam.ac.uk/kv227/papers/Discriminative_Training.pdf, 14 pages. cited by applicant

Xu et al. "Using social annotations to improve language model for information retrieval:" in Proceedings of the Sixteenth ACM Conference on Information and Knowledge Management, (CIKM '07). ACM, New York, NY, 2007, 1003-1006. cited by applicant

Zha et al. "Bipartite graph partitioning and data clustering:" in Proceedings of the Tenth

International Conference on Information and Knowledge Management (CIKM '01), Henrique Paques et al. (Eds.). 2001, ACM, New York, NY, 25-32. cited by applicant

Zweig et al. "Structured Models for Joint Decoding of Repeated Utterances", Interspeech 2008, 9th Annual Conference of the International Speech Communication Association, Brisbane, Australia, Sep. 22-26, 2008, 4 pages. cited by applicant

Zweig. "New Methods for the Analysis of Repeated Utterances", Interspeech 2009, 10th Annual Conference of the International Speech Communication Association, Brighton, United Kingdom, Sep. 6-10, 2009, 4 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/760,147, dated Nov. 23, 2011, 13 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,181, mailed Dec. 23, 2011, 23 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,181, dated Sep. 6, 2012, 26 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,172, mailed Dec. 29, 2011, 17 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,172, dated Jun. 29, 2012, 18 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/782,862, dated Nov. 30, 2011, 24 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/787,568, dated Nov. 22, 2011, 16 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,920, mailed Jan. 20, 2012, 17 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,175, mailed Jan. 10, 2012, 17 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,175, dated Aug. 13, 2012, 30 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,920, dated Nov. 2, 2012, 25 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/040,553, dated Jan. 10, 2012, 23 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/250,466, dated Dec. 29, 2011, 25 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/250,466, dated Jun. 5, 2012, 23 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/250,496, dated Dec. 13, 2011, 6 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/250,496, dated May 29, 2012, 18 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/077,106, dated Jun. 6, 2012, 20 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,180, mailed Dec. 27, 2011, 12 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,180, dated Jun. 22, 2012, 17 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/977,017, dated Sep. 26, 2012, 16 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,972, dated Jul. 13, 2012, 20 pages. cited by applicant

Authorized Officer H. Van Doremalen. International Search Report and Written Opinion in International Application No. PCT/US2012/021221, mailed Apr. 5, 2012, 11 pages. cited by applicant

Authorized Officer R. Poquet. International Search Report and Written Opinion in International Application No. PCT/US2011/037558, dated Jul. 29, 2011, 11 pages. cited by applicant

Authorized Officer R. Poquet. International Search Report and Written Opinion in International Application No. PCT/US2011/029407, dated Jun. 7, 2011, 10 pages. cited by applicant

Authorized Officer I. Barabasz. International Search Report and Written Opinion in International Application No. PCT/US2011/036984, dated Aug. 31, 2011, 13 pages. cited by applicant

Partial European Search Report in European Application No. 10165480.4, dated Apr. 28, 2011, 7 pages. cited by applicant

Extended European Search Report in European Application No. 10 165480.4, dated Oct. 24, 2011, 68 pages. cited by applicant

Chang et al. "A system for spoken query information retrieval on mobile devices." Speech and Audio Processing, IEEE Transactions on 10.8 (2002): 531-541. cited by applicant

Xu et al., "Cluster-based language models for distributed retrieval." Proceedings of the 22nd annual international ACM SIGIR conference on Research and development in information retrieval. ACM, 1999. cited by applicant

United States Office Action in U.S. Appl. No. 12/977,003, dated Mar. 27, 2014, 29 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,920, dated Apr. 24, 2014, 36 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/977,017, dated Apr. 15, 2013, 20 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/977,017, dated Oct. 7, 2013, 13 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,972, dated Mar. 5, 2013, 23 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,972, dated Nov. 5, 2013, 14 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,175, mailed Apr. 24, 2014, 29 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,180, dated Aug. 9, 2012, 22 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,181, dated Oct. 9, 2013, 37 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/977,017, dated Aug. 13, 2014, 21 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,972, dated Jul. 30, 2014, 18 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,181, dated Jul. 18, 2014, 30 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,180, mailed Nov. 28, 2014, 31 pages. cited by applicant

United States Office Action in U.S. Appl. No. 13/249,175, mailed Jan. 28, 2015, 30 pages. cited by applicant

United States Notice of Allowance in U.S. Appl. No. 12/977,003, dated Jan. 30, 2015, 9 pages. cited by applicant

United States Office Action in U.S. Appl. No. 12/976,920, dated Jan. 29, 2015, 31 pages. cited by applicant

United States Office Action in U.S. Appl. No. 14/299,837, dated Jan. 29, 2015, 30 pages. cited by applicant
Zha, Hongyuan, et al. "Bipartite graph partitioning and data clustering." Proceedings of the tenth international conference on Information and knowledge management. ACM, 2001. cited by applicant
United States Office Action in U.S. Appl. No. 12/976,972, dated Jul. 10, 2015, 15 pages. cited by applicant
United States Office Action in U.S. Appl. No. 12/976,920, dated Jul. 21, 2015, 20 pages. cited by applicant
United States Office Action in U.S. Appl. No. 13/249,175, mailed Jul. 16, 2015, 24 pages. cited by applicant
United States Office Action in U.S. Appl. No. 13/249,180, mailed Jul. 6, 2015, 28 pages. cited by applicant
United States Office Action in U.S. Appl. No. 14/299,837, dated Jul. 31, 2015, 8 pages. cited by applicant
United States Notice of Allowance in U.S. Appl. No. 14/299,837, mailed Oct. 13, 2015, 5 pages. cited by applicant
Office Action for the EP related application 16001249.8 dated Jun. 10, 2021. cited by applicant
USPTO. Office Action relating to U.S. Appl. No. 17/812,320, dated Oct. 5, 2023. cited by applicant

Primary Examiner: Chawan; Vijay B

Attorney, Agent or Firm: Honigman LLP

Background/Summary

CROSS REFERENCE TO RELATED APPLICATION (1) This patent application is a continuation of, and claims priority under 35 U.S.C. § 120 from, U.S. patent application Ser. No. 17/812,320, filed on Jul. 13, 2022, which is a continuation of U.S. patent application Ser. No. 16/892,749, filed on Jun. 4, 2020, which is a continuation of U.S. patent application Ser. No. 16/169,279, filed on Oct. 24, 2018, which is a continuation of U.S. patent application Ser. No. 14/988,408, filed on Jan. 5, 2016, which is a continuation of U.S. patent application Ser. No. 14/299,837, filed on Jun. 9, 2014, which is a continuation of U.S. patent application Ser. No. 13/249,172, filed on Sep. 29, 2011, which is a continuation of U.S. patent application Ser. No. 12/977,003, filed on Dec. 22, 2010, which claims priority under 35 U.S.C. § 119(e) from, U.S. Provisional Application 61/330,219, filed on Apr. 30, 2010, and U.S. Provisional Application 61/289,968, filed on Dec. 23, 2009. The disclosures of these prior applications are considered part of the disclosure of this application and are hereby incorporated by reference in their entireties.

TECHNICAL FIELD

(1) This document relates to systems and techniques for multi-modal input into an electronic device and conversion of spoken input to text.

BACKGROUND

(2) Computer users employ a number of mechanisms to provide input to their computing devices. Keyboards are common input devices, and they typically include single-digit numbers (e.g., in a cellular telephone) each of the letters in the alphabet, and some characters (e.g., in Qwerty or Dvorak keyboards). On mobile devices, keyboards are frequently "virtual" in form, and are displayed on a touch screen of a device. Such keyboards may be made available to various different applications running on a device, using a program known as an Input Method Editor, or IME, so

that the IME receives the user input and then passes it to whatever application is currently active on the device. An IME can also translate user input, such as when a user enters Roman characters in a written language like Pinyin, and the IME generates Chinese characters that correspond to the typed Pinyin. Where the Pinyin corresponds to multiple possible characters, the IME can display all such characters, the user can tap the intended character, and the IME can pass that character to the operating application.

(3) Users of computer devices, and particularly of mobile computing devices, may be constrained in their use of a keyboard. For example, the keyboard itself may be constrained in size because mobile device displays are small, so that only a sub-set of relevant characters can be displayed or the keys may be too small to press accurately. Also, the user may be constrained, in that they cannot easily type on a keyboard while walking through a crowded airport or driving a car. In such situations, spoken input may be preferred over typed input. However, speech-to-text conversion or translation typically requires lots of computer processing power, and mobile devices typically do not have much processing power. Also, such conversion often requires a particular user to “train” the system so that it better understands the user's voice and speech patterns.

SUMMARY

(4) This document describes systems and techniques for multi-modal input into an electronic device, including speech input. In one example, an IME that provides keyboard services to a device may also provide for spoken input to the device. Much of the IME services may be unchanged from an ordinary IME, but the IME may be augmented to provide for speech-to-text conversion.

Specifically, the IME may take a user input in whatever form (e.g., typed, spoken, D Pad, etc.) and may convert it to a standard form for presentation to an active application (e.g., passing characters or other key presses). Applications may subscribe to the IME's services just as they would subscribe if the IME did not accept voice input, and the provision of such speech-to-text functionality can be transparent to the applications.

(5) A user could choose when to provide typed input and when to provide spoken input to the IME. First, the user may be in an application that requires input, and may take an appropriate action to invoke the IME (e.g., pressing a particular button or soft key on the device, moving the device in a certain manner, and the like). A keyboard may be initially displayed on the user's device, and the user may take another appropriate action to indicate that he or she will provide speech input. For example, the user may press a microphone button on the virtual keyboard, or may make a swipe across the virtual keyboard or another action that is inconsistent with an intent to provide typed input in the virtual keyboard. At that point, the IME can begin “listening” to the device's microphone and after the user has spoken, may pass corresponding text to the application.

(6) In these manners, certain implementations may provide one or more benefits. For example, speech-to-text functionality may be provided on a computing device relatively simply, while re-using other IME functionality (e.g., interfaces to applications on a device) that is needed for keyboard-based IME translation. Use of context-specific language models in the manners discussed above and below may also permit more accurate conversion of speech to text, regardless of whether the system is trained to a particular user. Such context-specificity may also be provided automatically and transparently for a user, and at a level of specificity that is most suited to a particular situation. For example, where a user is simply interacting with a web page, the language model for the page may be used, but if the user is interacting with a form on the page, a more specific language model that is directed to that form or a particular field on the form may be used (with lower weightings, but higher-than-normal weightings, applied to the page-specific model).

(7) Particular manners of using public user activity, such as search activity, to build contextual language models may also result in the generation of accurate models in a convenient manner. Such data may be made available via natural user activity that is already occurring voluntarily by the users. Also, the activity may occur in large enough volumes to provide enough data needed to generate a complete and accurate model for many contexts, and in this case, for many web pages.

The models may also be updated over time, because the public user activity occurs continuously over time, so that continuously up-to-date models may be provided to users seeking speech-to-text conversion services.

(8) In general, in one aspect, methods, computer program products and systems are described for a multi-modal input-method editor. A request can be received from a user of an electronic device for an application-independent input method editor having written and spoken input capabilities. That the user intends to provide spoken input to the application-independent input method editor can be identified, and the spoken input can be received from the user. The spoken input can be input to an application executing on the electronic device. The spoken input can be provided to a remote server. The remote server includes a speech recognition system configured to recognize text based on the spoken input. Text can be received from the remote server, where the text represents the spoken input. The text can be provided to the application as user input.

(9) In general, in one aspect, methods, computer program products and systems are described relating to an input-method editor. A request is received from a user of an electronic device for an application-independent input method editor having written and spoken input capabilities. The application-independent input method editor is configured to receive input for multiple applications executable by the electronic device. It is identified that the user is about to provide spoken input to the application-independent input method editor. The spoken input is received from the user and corresponds to an input to an application from the multiple applications. The spoken input is provided to a remote server, which remote server includes a speech recognition system configured to recognize text based on the spoken input and is a server that is remote to the electronic device. Text is received from the remote server that represents the spoken input. The text is provided as the input to the application.

(10) Implementations of the methods, computer program products and systems can include one or more of the following features. A list of candidates of text representing the spoken input can be presented to the user. A selection can be received from the user of a candidate from the list. Providing the text to the application as user input can include providing the selection of the candidate from the list to the application.

(11) A written input can be received from the user and a language of the written input determined. A language indicator can be provided to the remote server based on the determined language indicating the language of the spoken input. A context indicator can be provided to the remote server such that the speech recognition system can select a language model from multiple language models based on the context indicator. The context indicator can specify the context in which the user input is received, a webpage in which the user input is received, an application in which the user input is received, a web form in which the user input is received, a field in the web form in which the user input is received and/or metadata associated with the field in the web form to name a few examples.

(12) A written input can be received from the user in a first writing system, and one or more candidates can be presented in a second writing system based on the written input in the first writing system. In an example, the first writing system is Pinyin and the second writing system is Hanzi.

(13) Intermediate processing can be performed on the received spoken input, e.g., noise reduction, filtering, or otherwise, and the spoken input provided to the remote server can be the intermediately processed spoken input.

(14) In some implementations, a request is received from a user of an electronic device for an application-independent input method editor having written and spoken input capabilities. The application-independent input method editor is configured to receive input for multiple applications executable by the electronic device. It is identified that the user is about to provide spoken input to the application-independent input method editor. The spoken input is received from the user and corresponds to an input to an application from the multiple applications. The spoken input is

converted at the electronic device to text that represents the spoken input. The text is provided as the input to the application.

(15) The details of one or more embodiments are set forth in the accompanying drawings and the description below. Other features and advantages will be apparent from the description and drawings, and from the claims.

Description

DESCRIPTION OF DRAWINGS

(1) FIG. 1A is a conceptual diagram of an example system including a multi-modal input method editor.

(2) FIG. 1B is a block diagram of an example system that can be used to implement the multi-modal input method editor.

(3) FIG. 1C is a block diagram of example software that can be used to implement the input method editor.

(4) FIG. 2 is a flow chart of an example process for generating text using a multi-modal input method editor.

(5) FIGS. 3A-E show example screen shots of an electronic device including a multi-modal input method editor.

(6) FIG. 4 is a block diagram of an example system for receiving speech input and training language models to interpret the speech input.

(7) FIG. 5A is a block diagram of queries and associated websites in search results

(8) FIG. 5B is a URL showing search terms in a query

(9) FIG. 6 is a block diagram of an example system for recognizing speech input to a computing device.

(10) FIGS. 7A and 7B show flowcharts of example processes of building an interpolated language model.

(11) FIG. 8 shows a flowchart of an example process for recognizing text in an utterance.

(12) FIG. 9 shows examples of generic computer devices that may be used to execute the actions discussed in this document.

(13) FIG. 10 is a block diagram of an example data structure of a language model.

(14) Like reference symbols in the various drawings indicate like elements.

DETAILED DESCRIPTION

(15) This document describes techniques and systems that may be used to provide speech-to-text conversion for a user of a computing device, such as a smartphone. In certain instances, the speech input may be handled in a manner similar to other input (e.g., typed input) using an application such as an IME, where the IME can be switched into modes depending on the manner that the user chooses to enter data. Where transformation of the input is needed, the input may in certain circumstances be transmitted (either in a raw or converted form) to a server system remote from the computing device that may be programmed to pass a transformed input back to the device, such as by providing text in response to receiving spoken inputs. The computing device may also provide the server system with meta data that is passed with, or at essentially the same time as, the spoken inputs, and the meta data may be used by the server system to identify a context in which the user is entering the spoken input. The server system may then use that meta data to identify a language model to be used and/or to build a language model on the fly, such as by dynamically applying particular weightings to different language models (which may each be derived from different input corpuses).

(16) FIG. 1A is a conceptual diagram of an example system **100** that includes a multi-modal input method editor (IME) **102**. In this example, the IME **102** is implemented in a mobile electronic

device **104**, though it should be understood that the IME **102** can be implemented in a different electronic device, e.g., a PC, laptop computer, PDA, etc. The electronic device **104** includes multiple user input devices **106**, including a microphone **105** to receive spoken user input. Other user input mechanisms include a keyboard, which can include a soft or virtual keyboard (e.g., a touchscreen keyboard **108**) or a hard or physical keyboard, a mouse, a trackball, and the like. The user input mechanisms **106** are capable of receiving spoken input (i.e., by the microphone) and written input (i.e., by the keyboard **108**).

(17) The user input can be received by the electronic device **104** for use as input into one of various applications **110** that can execute on the electronic device **104**, e.g., a web browser, an e-mail application, a word processing application, a contacts book, and/or a calendar. In some implementations, the user input is an input into a web form on a particular web page of a particular web site. The IME is generally application-independent, i.e., can be used for most any of the applications **110**.

(18) If the user input is spoken input, i.e., a speech utterance, the spoken input can be provided to a remote server for conversion to text. For example, the speech utterance **111** can be transmitted over the network **112** to a remote server **114** that includes a speech service **116** and speech recognizer system **118**. The network **112** can include one or more local area networks (LANs), a wide area network (WAN), such as the Internet, a wireless network, such as a cellular network, or a combination of all of the above.

(19) The speech recognizer system **118** can use one or more language models **120** to recognize text from the speech utterance. The text **113**, which can be a selected best candidate or can be a list of n-best candidates that correspond to the speech utterance, is provided back to the electronic device **104** over the network **112**. The text **113** can be displayed to the user on a display **122** of the electronic device **104**.

(20) If the text **113** includes a list of n-best candidates, the user **101** can select a candidate from the list that corresponds to the user's spoken input, for example, using the keyboard **108** or another input mechanism, such as touching the touch screen over one of the candidates, to navigate the list and make a selection.

(21) The user can also provide written input, and can provide input using a combination of written and spoken input. For example, the user can begin a search query in a web browser by speaking one or more words and can then add to the query string by typing additional input using the keyboard **108**. The IME **102** can provide the combined user input to the relevant application, i.e., the web browser application in this example. In some implementations, the language that the written input is written in can be determined and then provided as a language indicator to the remote server **114**. The remote server **114** can use the language indicator when converting the speech utterance **111** to the text **113**. For example, by knowing the language in which the speech is spoken, an appropriate language model **120** can be selected for use by the speech recognizer **118**.

(22) FIG. 1B is a block diagram of an example system **130** that can be used to implement a multi-modal IME. The example system **130** can be implemented, for example, in a computer device, such as a personal computer device, or other electronic devices, such as a mobile phone, mobile communication device, personal digital assistant (PDA), Global Positioning System (GPS) navigation device, and the like.

(23) The example system **130** includes a processing device **132**, a first data store **134**, a second data store **136**, input devices **138**, output devices **140**, and a network interface **142**. A bus system **144**, including, for example, a data bus and a motherboard, can be used to establish and control data communication between the components **132**, **134**, **136**, **138**, **140** and **142**. Other system architectures can also be used.

(24) The processing device **132** can, for example, include one or more microprocessors. The first data store **134** can, for example, include a random access memory storage device, such as a dynamic random access memory, or other types of computer-readable medium memory devices.

The second data store **136** can, for example, include one or more hard drives, a flash memory, and/or a read only memory, or other types of computer-readable medium memory devices.

(25) The input devices **138** include at least one input device that is configured to receive spoken input and at least one input device configured to receive written input. Example input devices **138** can include a microphone, keyboard, a mouse, a stylus, etc., and examples **140** can include a display device, an audio device, etc. The network interface **142** can, for example, include a wired or wireless network device operable to communicate data to and from a network **146**. The network **146** can include one or more local area networks (LANs), a wide area network (WAN), such as the Internet, a wireless network, such as a cellular network, or a combination of all of the above.

(26) In some implementations, the system **130** can include input method editor (IME) code **131** from a data store, such as the data store **136**. The input method editor code **131** can be defined by instructions that upon execution cause the processing device **132** to carry out input method editing functions. The input method editor code **131** can, for example, include interpreted instructions, such as script instructions, e.g., JavaScript or ECMAScript instructions, that can be executed in a web browser environment. Other implementations can also be used, e.g., a stand-alone application, an applet, a plug-in module, etc., for use in a user interface, such as a display that displays user inputs received by use of keypad mapping for a mobile device or keyboard mapping for a mobile device or personal computer.

(27) Execution of the input method editor code **131** generates or launches an input method editor instance (IMEI) **133**. The input method editor instance **133** facilitates the processing of one or more input methods at the system **130**, during which time the system **130** can receive inputs for characters or symbols, such as, for example, spoken or written input. For example, the user can use one or more of the input devices **138**, e.g., a microphone for spoken input or a keyboard for written input. In some implementations, the user input can be Roman characters that represent input in a first writing system, e.g., Pinyin, and the input method editor can convert the input to a second writing system, e.g., Hanzi terms. In some examples, a Hanzi term can be composed of more than one Pinyin input.

(28) The first data store **134** and/or the second data store **136** can store an association of inputs. Based on a user input, the input method editor instance **133** can use information in the data store **134** and/or the data store **136** to identify one or more candidate selections represented by the input. In some implementations, if more than one candidate selection is identified, the candidate selections are displayed on an output device **140**. For example, if the user input is spoken input, then a list of candidate selections showing written text representations of the spoken input can be presented to the user on the output device **140**. In another example, if the user input is Pinyin inputs, the user can select from the candidate selections a Hanzi term, for example, that the user desires to input.

(29) In some implementations, a remote computing system **148** having access to the system **130** can be used to convert spoken user input to written user input. For example, the remote system **148** can be a server that provides a speech recognition service via the network **146**. One or more speech utterances forming the spoken input can be transmitted to the remote system **148** over the network **146**. The remote system **148** can determine a text conversion of the spoken input, for example, using a convenient form of speech recognizer system, and transmit the text conversion to the system **130**. The text conversion can be a best candidate for text corresponding to the spoken input or can be a list of n-best candidate selections for presentation to the user for selection as the input. In an example implementation, the speech recognizer system can include Hidden Markov Modeling (HMM) encoded in a finite state transducer (FST). Other configurations of speech recognizer can be used by the remote system **148**.

(30) In some implementations, the remote system **148** can also be used to edit a logographic script. For example, the remote system **148** may be a server that provides logographic script editing capability via the network **146**. In one example, a user can edit a logographic script stored in the

data store **134** and/or the data store **136** using a remote computing system, e.g., a client computer. The system **130** can, for example, select a character and receive an input from a user over the network interface **142**. The processing device **132** can, for example, identify one or more characters adjacent to the selected character, and identify one or more candidate selections based on the received input and the adjacent characters. The system **130** can transmit a data communication that includes the candidate selections back to the remote computing system.

(31) FIG. **1C** includes a block diagram of example software that can be used to implement an input method editor in FIG. **1B** (e.g., IMEI **133**). The system **160** includes a user interface **162** and software **164**. A user **166** can access system **160** through the user interface **162**. The software **164** includes applications **165**, IME engine **166**, an operating system (OS) **167**, a speech recognition system **169** including a language model **168**, and a detection engine **170**. The operating system **167** is a particular piece of software that can provide the user interface **162** between the software **164** (e.g., applications **165** and IME engine **166**) and the user **166**.

(32) As shown in FIG. **1C**, the speech recognition system **169** and language model **168** are separate from IME engine **166**. In particular, the speech recognition system **169** and language model **168** (which can include two or more language models) are included within software **164** as a separate software component. Other implementations are possible. For example, the speech recognition system **169** and language model **168** can be located remotely (e.g., at the remote system **148** of FIG. **1B**). As another example, the speech recognition system **169** and language model **168** can be included within the IME engine **166**.

(33) The language model **168** can define one or more language sub-models, each sub-model tailored to a particular application, or webpage, or webform on a particular webpage, or website, to name a few examples. Each language sub-model can, for example, define a particular rule set, e.g., grammar particular to a language, phrase sets, verbals, etc., that can be used to determine a user's likely intent in entering a set of inputs (e.g., inputs for generating candidates that are translations, transliterations, or other types of phonetic representations). In some implementations, each language sub-model can also include a user history of a particular user, e.g., a dictionary of words and phrased often used by a particular user.

(34) The detection engine **170** includes an input module **172** and can include a timing module **174**. The input module **172** can, for example, receive input (e.g., keystrokes representing characters or a speech utterance) to particular applications **165** and send the received input to the IME engine **166**. In some implementations, the detection engine **170** is a component of the IME engine **166**.

(35) The detection engine **170** can detect input and determine whether or not to send the input to the IME engine **166**. The IME engine **166** can, for example, be implemented using the input method editor code **131** and associated data stores **134** and **136**, and provide output candidates in text converted from speech to an interface (e.g., user interface **162**) as the input (e.g., speech utterances) is detected, as described with reference to FIGS. **2** and **3A-E** below.

(36) The components of system **160** can be communicatively coupled to one or more of each other. Though the components identified above are described as being separate or distinct from each other, one or more of the components may be combined in a single system, or to perform a single process or routine. The functional description provided herein including separation of responsibility for distinct functions is by way of example. Other storage architectures can also be used. In particular, other groupings or other divisions of functional responsibilities can be made as necessary or in accordance with design preferences. For example, IME engine **166** can perform the functions of detection engine **170**. As another example, input module **172** and timing module **174** can be combined into a single module.

(37) FIG. **2** is a flowchart of an example process **200** for using an input method editor to receive spoken input from a user input device and to provide written, or textual, input to a corresponding application. A request is received from a user for an application-independent input method editor that has written and spoken input capabilities (Step **202**). By way of illustrative example, and

without limitation, the request can be received by a mobile electronic device that has a touchscreen keyboard. Example screenshots from such a mobile electronic device are shown in FIGS. 3A-E. These screenshots can be used to illustrate the example process **200**; however, it should be understood that other devices can implement the process **200**, and the screenshots shown are not intended to be limiting.

(38) FIG. 3A shows a screenshot **300** where a user has selected to activate a web browser application. Through the web browser application, the user has selected to navigate to the Google search page at the URL www.google.com **302**. FIG. 3B shows a screen shot **304** with a soft touchscreen keyboard **306** displayed in a lower portion of the display screen. For example, the user can touch or tap the screen in the search query field **308** to automatically have the keyboard **306** displayed, although other mechanisms can be used to trigger the display of the keyboard **306**. The example keyboard **306** shown includes a microphone key **310**. An example of a request that can be received from the user includes the user selecting the microphone key **310**. Another example includes the user selecting a graphical entity, such as a microphone icon or button, displayed next to or in an input field, e.g., in search query field **308**. Another example includes the user swiping his/her finger across the input field, e.g., in a left to right motion, or tapping the input field. Yet another example includes the user picking up the device in a manner that is consistent with raising a microphone included in the device to the proximity of the user's mouth, which can be detected, for example, by an accelerometer reading. Other forms of request can be received from the user for an application-independent input method editor having written and spoken input capabilities, and the above are but some examples.

(39) A user's intention to provide spoken input to the application-independent input method editor is then identified in the process (Step **204**). For example, receiving a speech utterance from the user can be used to identify that the user intends to provide spoken input. In other implementations, receiving the request from the user for the input method editor with written and spoken input capabilities can also be used to identify that the user intends to provide spoken input, i.e., the same user action can provide both the request and be used to identify the user's intention. In some implementations, as shown in the screenshot **312** in FIG. 3C, a graphical element can be displayed that prompts the user to speak, such as the microphone graphic **314** and the instructions "Speak now" **316**.

(40) A spoken input, i.e., a speech utterance, is then received from the user. The user provides the spoken input as input to an application that is executing on the device (Step **206**). The spoken input is provided to a remote server that includes a speech recognition system configured to recognize text based on the spoken input (Step **208**). For example, referring again to FIG. 1B, the spoken input can be sent over the network **146** to the remote system **148**, where the remote system **148** includes a speech recognition system to recognize text from a speech utterance. Because processing the speech to text conversion can take some time, in some implementations a graphic is displayed to the user to indicate that the process is in progress, such as the "Working" graphic **320** shown in the screenshot **322** in FIG. 3D.

(41) Text is then received from the remote server, where the text represents the spoken input (Step **210**). Once the remote server, e.g., remote system **148**, has processed the speech utterance, the corresponding text is sent back to the user's device and can be displayed for the user. In some implementations, the best candidate for representation of the speech utterance is selected by the speech recognition system at the remote server and provided to the device. However, in some implementations, an n-best list of candidates can be provided and presented to the user for selection of the correct candidate. For example, referring to FIG. 3E, a screen shot **324** shows a list of suggestions, with the best candidate "the man in the moon" displayed at the top of the list as the default selection.

(42) The text, i.e., the spoken input converted to written input, is then provided to the application as user input (Step **212**). That is, once the correct text conversion is selected, if a list of candidates

was provided, or once the best candidate has been received, if only one was sent from the remote server, the written input can be passed to the application as the user input for processing by the application.

(43) In some implementations, a context indicator can be sent with the spoken input to the remote system for conversion to text. The remote system can use the context indicator to facilitate the speech-to-text conversion. For example, the context indicator can be used as a basis for selecting an appropriate language model to use by the speech recognition system. The context indicator can specify the context in which the spoken user input was received. For example, the context indicator can specify a name of a field, e.g., in a web form, the name of the application in which the input was received, and/or identify a web page if the user input was received in a web browser application. As another example, the context indicator can include metadata relating to a field in which the user input was received. For example, the metadata can specify that the field requires a one-word answer, or a date, or a name, and the like. In some implementations, the context indicator information can be obtained by the input method editor from the operating system of the electronic device.

(44) The device may pass a sound file (including streaming sound data) of the spoken input to a remote server system, and may take steps to improve the quality of the speech-to-text conversion. As one example, the device may pass information that allows the server system to select a language model that is relatively small in size and is specific to the task that the user is currently facing. For example, when applications register with the IME, they may provide information about fields into which a user can enter information in the applications. The IME can pass such information to the server system, so that the server system may select an appropriate language model. For example, if the cursor is in an “address” field of an application, the IME can pass such information to the server system so that, for example, a user utterance that sounds like “scheet” is interpreted as “street” and not “sweet.”

(45) The language model that is applied may also be a composite of, or interpolation of, multiple separate language models. The different models may be relatively small models that have been derived from large data corpuses (e.g., SMS messages and e-mail messages). The models may be prepared off-line by analyzing such corpuses, and the mixture weights that are applied to the models may be generated on-the-fly at run-time, including after data from a particular instance of speech input starts being received from a user. The weightings may be a function, for example, of the field into which a user is currently making an utterance (e.g., the “to” or “from” fields of an email message versus the “body” field of an email message).

(46) The language model may also be chosen more generally, without reference to a particular input field in which the cursor is located. For example, an application developer may register their application with an app store, and may indicate the type of application that it is, such as a music player. Similarly, a server system may have a number of topic-specific language models that it stores. Thus, if a user of a music application speaks the word “Heart” or “Hart” (which would not be plain from the spoken word itself), the IME may pass a “music” indication to the server system, so that the application is passed the word “heart,” and the user sees a song list for the female-led rock band. If the media player is a video player, the IME may pass the word “Hart” (assuming there are no great movies or television shows with the word “heart” in their titles) so that the user is shown an index of the episodes of the iconic detective drama “Hart to Hart.”

(47) When the user is interacting with a web page, such as in a web browser, different context information can be used for selecting the proper language model. In particular, the domain or a sub-domain for the page may be provided, so that the that the language model will be specific to the particular type of web site. For example, if the web site is Amazon, then the language model may be one in which “shopping” terms have higher prominence. For example, “product” may have a higher score than “protect” for similar sounds. Such a model may be prepared to be directed to the site itself (e.g., by analyzing input forms on the site, and analyzing text on the site), or on a

category that the site matches. Thus, for example, the same language model may be used for the sites Amazon.com, Buy.com, and the like.

(48) FIG. 4 is a block diagram of an example system **400** for receiving speech input and training language models to interpret the speech input. In the system **400**, users of client devices enter data into text input fields, and a speech server analyzes that data to determine the type of text the users entered. The speech server builds language models from this information, and uses the language models to recognize speech input associated with similar input fields.

(49) Clients **402**, such as computers executing a web browser with an optional toolbar, can receive text input for web forms. When this input is used as part of a process to browse to another page, such as typing in a search field or order form field before pressing a submit button, the text and destination information can be saved in a toolbar log **404**. Clients **406**, such as computers and mobile computing devices, can receive text based web search queries from users. These queries can be resolved by a web search server (not shown) and the queries and search results can be stored in query logs **408**. The client **406** can receive text input to applications, such as an email client, a messaging client, and/or a word processor and spoken input to applications. The text input and transcriptions of spoken input can be stored in input logs **409**.

(50) A speech recognition server **410** can retrieve the data in the toolbar logs **404**, the query logs **408**, and/or the input logs **409**. A speech recognition system **412** can group this data into categories or classifications. The speech recognition system **412** creates a series of language models **414a-414n**. The language models can contain words, phrases, sentences, etc. from the toolbar logs **404**, query logs **408**, and/or the input logs **409** based on a particular topic. For example, if the language model A **414a** is a language model of British last name, it can contain “Churchill,” “Bean,” and “Pigou.” In another example, if the language model B **414b** contains technology-related language, it can contain “object oriented programming,” “reverse polish notation,” and “garbage in, garbage out.”

(51) In some implementations, the speech recognition system **412** can build interpolated language models from the language models **414a-414n**. For example, the speech recognition system **412** can use an utterance with a known text result, and poll the language models **414a-414n** to determine the confidence level that each language model **414a-414n** would associate with the utterance/known text pair. The confidence level returned by each language model **414a-414n** can be used by the speech recognition system **412** to create an interpolated language model. For example, using a test utterance and text of “Mr. Bean writes in reverse polish notation,” the language model A **414a** gives a confidence level of 50%, the language model B **414b** give a confidence level of 65%, and the other language models **414c-414n** give substantially 0% confidence levels. In this example, the speech recognition system **412** can create an interpolated language model that heavily weighs the language model A and B **414a-b**. The particular weightings may be equal to or proportional to the confidence levels, in certain implementations.

(52) A client device **416** can execute an application that accepts text input and can receive an utterance from a user for that text input. The client device **416** can transmit the utterance, along with the context in which the utterance was made or received, to the speech recognition server **410** for translation into text. In some implementations, the context in which an utterance is received is determined by the client device **416** with the field name or label associated with the text input. For example, some applications include metadata for an application input field such as package name, field number or name, and/or attribute flags (e.g. long sentence, email, street address, etc.). In some applications, a text label associated with or displayed near an input box is used to determine the context.

(53) A speech service **418** can receive the utterance and context. In some implementations, the speech service **418** can map the context to broader categories or to the categories of the language models **414a-414n**. In some implementations, the speech service **418** can maintain a cache of utterances and resultant texts, and if a received utterance substantially matches a cached utterance,

the speech service **418** can return the appropriate cached resultant texts.

(54) The speech recognition system **412** uses or creates an interpolated language model to recognize the text in the utterance. In some implementations, a single candidate text is identified. In some implementations, the top n candidates are identified either wherein all candidates meet a certain confidence threshold, or wherein the top n candidates are selected.

(55) The text or texts identified by the speech recognition system **412** is returned to the client **416**, where it is, for example, displayed in a text input field that has focus.

(56) Such context-specific language models may be generated in a variety of ways, including by using search information from a public search engine. For example, bi-grams may be made between search queries and web sites that are highly relevant to those queries (e.g., are in the list of top **10** sites returned to the queries). The queries and sites may then be assigned as nodes in a graph, joined by the relationship between the two. Clusters of such relationships may then be identified, and language models may be generated for those clusters. A particular language model for a web site may then be generated as a weighted combination of the language models for the clusters around the site in the graph.

(57) FIG. 5A is a block diagram **500** of queries and associated websites in search results. A group of queries **502** is a collection of example queries that can be collected from a browser executing a toolbar. The queries can be submitted to a web search engine, and a list of search results can be returned. For example, a search for “mike lebeau” can result in search results including facebook.com and linkedin.com. A search for “blink” can result in search results including amazon.com.

(58) The queries and associated websites can be used, for example, in determining a relationship between queries and related topics or categories. In this example, if it is known that facebook.com and linkedin.com are social network websites, it can be determined that the queries “bill byrne,” “will rusch,” “mike lebeau,” and “brian stope” may be the names of people. Similarly, if it is known that amazon.com is a retailer with a reputation as a book retailer, it can be determined that “blink” and “great gatsby” may be the names or titles of a retail products, perhaps books. Similarly, if a search for “william gibson” returns links to both facebook.com and amazon.com, it can be determined that “william gibson” may be an author.

(59) FIG. 5B is a URL showing search terms in a query. The URL in this example can be recorded by a web browser executing a toolbar. When a user submits a form via a GET request, the contents of the form can be encoded and logged in the URL parameters. In this example, a search on the webpage facebook.com can generate the URL shown. The name parameter **550**, in this example “bill %20byrne”, indicates that “bill byrne” was entered into a field in a web form.

(60) FIG. 6 is a block diagram of an example system **600** for recognizing speech input to a computing device. In the system **600**, a user enters speech input to an application via an IME. The speech input is analyzed at a speech recognition server and text is returned to the application.

(61) An application **604** executes on a computing device **602**, for example a mobile computing device. An input method manager **608** can manage input methods that the application **604** accepts, including speech input from a user. A speech IME **610** in the operating system of the computing device **602** records speech input and collects context data from an input context **612**.

(62) Speech IME **160** can transmit the speech input and context to a voice data center **614**. A speech service front end **616** can receive the speech input and context and prepare the speech input and context for analysis. For example, statistics can be collected, recognition jobs can be created and sorted, etc. A speech recognition system **618** examines the context, and selects an associated language model from a big table of language models **622**. The associated language model in the big table of language models **622** can contain a reference to a language model in a ProdLM **620**. The speech recognition system **618** uses the referenced language model in the ProdLM **620** to analyze the speech input and determine a text string. The text string is returned to the speech input method **610**, for example by way of the speech service front end **616**. The text string can be presented to

the user of the computing device **602**, and if approved by the user, can be sent to the application **604** as text input.

(63) In some implementations, data obtained from web toolbar logs can be used in determining an interpolated language model. A sample of content of user input (e.g., queries) into a web form can be obtained from web toolbar logs, e.g., in URLs. That is, when a user submits a web form via a GET request, the content of the user input into fields of the web form can be encoded and logged in the URL parameters (e.g., as was described above in reference to FIG. 5B). A toolbar log can therefore provide sample data of what users have input (i.e., typed) into each of the fields of a number of web forms. In some instances, however, the volume of this data may be too sparse to train a good language model for the web form. However, for a particular web form a larger training set can be available if queries of the same category as the queries input into the particular web form are used. That is, queries that correspond to search results having similar websites in the top N search results (where N is an integer, e.g., **10**) as each other are probably generally within the same category. To group queries by category, clusters in a search query—search result (website) graph (where queries and websites are assigned to nodes) can be identified, where an edge between a particular query node and a particular website node occurs whenever the website appears in the top N results for the particular query. A language model can then be trained per cluster, that is, language models can be trained for a category (if a cluster is representative of a category of queries).

(64) To generate an interpolated language model for a particular web form (or web site), the interpolated language model can be built using multiple language models formed from a number of clusters. In some implementations, the multiple language models can be selected as follows. A sample of user input into the particular web form (i.e., sample queries) can be obtained, e.g., from a web toolbar log. For example, referring again to FIG. 5B, the sample data may be the query “bill byrne” that is input into a web form on the www.facebook.com website, which is also reflected in the graph shown in FIG. 5A. In this example, “bill byrne” belongs to a cluster that includes people-oriented websites (e.g., facebook.com and linkedin.com) and people names (Bill Byrne, Will Rusch, Mike Lebeau, Brian Strobe and William Gibson). In this simple example, the “bill byrne” query belongs to only the people names cluster. However, a query may belong to many clusters. In such instances, for the sample queries obtained from the web toolbar data, the K most significant clusters (K is a number/integer) can be identified and weighted, with the sum of their weights adding to 1.0. In some implementations, the K most significant clusters can be determined by counting the number of sample queries (from the web toolbar log sample data) that fall into each cluster and then optimizing the weights of only the top K clusters. Other techniques to determine the K most significant clusters can be used. The interpolated language model for the web form can be used to convert a speech input received for a field in the form into text. The converted text can be transmitted to an electronic device that initially received the speech input, or transmitted elsewhere, to be used as text input for the field in the web form.

(65) FIG. 7A shows a flowchart of an example process **700** of building an interpolated language model for a website (or a web form of a website). In the process **700**, pairs of queries and results are collected, base language models are created, and interpolated language models are created.

(66) In step **702**, pairs containing web queries and result sites are extracted from web search logs. For example, a web browser or web browser add on can report queries, query result URLs, or search results to a central repository. In another example, a web search server can create a log of search terms and website domain names that were returned in response to those search terms.

(67) In step **704**, a cluster bipartite graph, also known as a bigraph, is formed by the pairs. A bipartite graph is a graph whose vertices can be divided into two disjoint sets such that every edge connects a vertex in the first set to one in the second set. In some implementations, the first set can contain nodes representing queries and the second set can contain nodes representing results. Clusters are identified in the bipartite graph, for example so that each cluster consists of queries

and results that may be by semantically or syntactically similar.

(68) In step **706**, language models are trained based on the queries in the clusters. In some implementations, sample utterances are recorded by a range of speakers and associated with the queries. The utterances for each query can be aggregated, such as by determining an average waveform and a deviation factor indicating how much deviation from the mean was found for most utterances.

(69) In step **708**, web form sample data is obtained from toolbar logs, the web form sample data relating to the particular web form (or website) for which the interpolated language model is being generated (e.g., the interpolated language model can be for the web form). In some implementations, a web browser toolbar can record and transmit the URLs of websites visited by the browser. The URLs can include input text (i.e., sample queries) that has been entered into the web form and submitted. This information can be used to associate a particular website domain, or web page with a particular topic or idea.

(70) In step **710**, the K most significant clusters for the input text, i.e., the sample queries, obtained from the web form sample data are identified. A list of cluster names is collected and ranked in order of significance or similarity to the data collected in the step **708**. A predetermined number (K, in this case) of the most significant clusters are identified. For example, clusters are identified that include the sample queries. To limit the number of clusters to K, i.e., to identify the K most significant, in some implementations, the number of sample queries that fall into each cluster can be counted and the weights assigned to each cluster optimized for only the top K clusters.

(71) In step **712**, an interpolated language model for the web form (or website) is built. In one example, a data structure is created that contains links to the K most significant language models. This data structure represents a combined language model that delegates speech recognition functionality to the referenced language models.

(72) In step **714**, K interpolation weights are optimized. Each referenced language model can have a weight, such as a percentage, integer in a particular range, or probability, associated with the referenced language model in the interpolated language model. In some implementations, this weight can be determined by testing a known pairs of utterances and text. The pairs can be submitted to the referenced language models, which can return a confidence level representing the likelihood or accuracy of the language model to correctly identify the text in the utterance. The confidence levels, optionally normalized or otherwise converted, can be used as weighting values in the interpolated language model. Thus, reference language models likely to correctly identify text in a particular classification are most heavily weighted.

(73) Although a particular number, type, and order of steps are shown, it will be understood by one skilled in the art that other number, types, and orders are possible. For example, other methods of clustering or graph creation can be used in the step **704**. In another example, input field flags and descriptions are obtained in the step **708**. In another example, step **702** and step **714** can be combined using a process that builds an interpolation model and interpolation weights in one step.

(74) FIG. 7B shows a flowchart of an example process **750** of building an interpolated language model. In the process **750**, text input and associated metadata is collected, base language models are created, and interpolated language models are created using the base models and interpolation factors such as weightings to be provided to each of multiple base models.

(75) In step **752**, text input and input metadata pairs are extracted from input logs. For example, an application can report to a central repository input entered into an input field and metadata for the input field. In some examples, the text input can include voice input that is transcribed into a text format. In some examples, application wide or operating system wide metadata schemes can define or describe the type of input field, such as an address field, a free form text field, a search field, or a social status field.

(76) In step **754**, a cluster bipartite graph, also known as a bigraph, is formed by the pairs. A bipartite graph is a graph whose vertices can be divided into two disjoint sets such that every edge

connects a vertex in the first set to one in the second set. In some implementations, the first set can contain nodes representing input metadata and the second set can contain nodes representing text input. Clusters are identified in the bipartite graph, for example so that each cluster consists of similar input metadata.

(77) In step **756**, language models are trained based on the text input in the clusters. In some implementations, sample utterances are recorded by a range of speakers and associated with the metadata. The utterances for each text input can be aggregated, such as by determining an average waveform and a deviation factor indicating how much deviation from the mean was found for most utterances.

(78) In step **758**, input metadata is categorized. The input metadata can be used to associate a particular input field or application with a particular topic or idea. Example topics or ideas include, but are not limited to, address fields, free form text fields, search field, social status fields, and numeric fields.

(79) In step **760**, the K most significant clusters for a category are identified. A list of cluster names is collected and ranked in order of significance or similarity to the data collected in the step **758**. A predetermined number (K, in this case) of the most significant clusters are then identified.

(80) In step **762**, an interpolated language model is built. In one example, a data structure is created that contains links to the K most significant language models. This data structure represents a combined language model that delegates speech recognition functionality to the referenced language models.

(81) In step **764**, K interpolation weights are optimized. Each referenced language model that is referenced in the data structure can have a weight, such as a percentage, integer in a particular range, or probability, associated with it. In some implementations, this weight can be determined by testing a known pair or pairs of utterances and text. The pairs can be submitted to the referenced language models, which can return a confidence level representing the likelihood or accuracy of the language model to correctly identify the text in the utterance. The confidence levels, optionally normalized or otherwise converted, can be used as weighting values in the interpolated language model. Thus, reference language models that are likely to identify text in a particular classification correctly are most heavily weighted.

(82) Although a particular number, type, and order of steps are shown for the process in this figure, it will be understood by one skilled in the art that other number, types, and orders are possible. For example, in step **762** and step **764**, a data structure with weighted links to the K most significant language models can be selected from a collection of some or all possible data structures with weighted links to language models.

(83) FIG. **8** shows a flowchart of an example process **800** for recognizing text in an utterance. In the process **800**, an utterance is received, the speech in the utterance is recognized, and text, or a list of text options, is returned.

(84) In step **802**, an utterance is received. For example, a speech recognition server can receive translation requests from clients. The translation request can include an utterance (e.g. an audio file) and a context (e.g. text or other data describing how the utterance may be used or categorized).

(85) In step **804**, a relevant language model is determined. For example, an interpolated language model for web search, free form text input, or social status can be determined. In some implementations, language models, including interpolated language models, can be indexed by keyword, web domain, application type, or other criteria. Using metadata associated with the utterance, such as context or source information, a relevant language model can be determined.

(86) In step **806**, speech recognition is performed. The utterance can be used by a speech recognition application using the language model. The speech recognition application can calculate one or more candidate text strings from the utterance. A confidence level can be associated with each text string.

(87) In step **808**, a K-best list of text strings representing the utterance is returned (where K is an

integer). In some implementations, the text string with the highest confidence value is selected and returned. In some implementations, a particular number (K, in this case) of text strings with the highest confidence level are returned.

(88) Although a particular number, type, and order of steps are shown, it will be understood by one skilled in the art that other number, types, and orders are possible. For example, in step **808**, all text strings with a confidence level above a particular threshold can be returned. In another example, an utterance can be preprocessed before step **806** to improve recognition.

(89) FIG. **9** shows an example of a generic computer device **900** and a generic mobile computer device **950**, which may be used with the techniques described here. Computing device **900** is intended to represent various forms of digital computers, such as laptops, desktops, workstations, personal digital assistants, servers, blade servers, mainframes, and other appropriate computers. Computing device **950** is intended to represent various forms of mobile devices, such as personal digital assistants, cellular telephones, smartphones, and other similar computing devices. The components shown here, their connections and relationships, and their functions, are meant to be exemplary only, and are not meant to limit implementations of the inventions described and/or claimed in this document.

(90) Computing device **900** includes a processor **902**, memory **904**, a storage device **906**, a high-speed interface **908** connecting to memory **904** and high-speed expansion ports **910**, and a low speed interface **912** connecting to low speed bus **914** and storage device **906**. Each of the components **902**, **904**, **906**, **908**, **910**, and **912**, are interconnected using various busses, and may be mounted on a common motherboard or in other manners as appropriate. The processor **902** may process instructions for execution within the computing device **900**, including instructions stored in the memory **904** or on the storage device **906** to display graphical information for a GUI on an external input/output device, such as display **916** coupled to high speed interface **908**. In other implementations, multiple processors and/or multiple buses may be used, as appropriate, along with multiple memories and types of memory. Also, multiple computing devices **900** may be connected, with each device providing portions of the necessary operations (e.g., as a server bank, a group of blade servers, or a multi-processor system).

(91) The memory **904** stores information within the computing device **900**. In one implementation, the memory **904** is a volatile memory unit or units. In another implementation, the memory **904** is a non-volatile memory unit or units. The memory **904** may also be another form of computer-readable medium, such as a magnetic or optical disk.

(92) The storage device **906** is capable of providing mass storage for the computing device **900**. In one implementation, the storage device **906** may be or contain a computer-readable medium, such as a floppy disk device, a hard disk device, an optical disk device, or a tape device, a flash memory or other similar solid state memory device, or an array of devices, including devices in a storage area network or other configurations. A computer program product may be tangibly embodied in an information carrier. The computer program product may also contain instructions that, when executed, perform one or more methods, such as those described above. The information carrier is a computer- or machine-readable medium, such as the memory **904**, the storage device **906**, memory on processor **902**, or a propagated signal.

(93) The high speed controller **908** manages bandwidth-intensive operations for the computing device **900**, while the low speed controller **912** manages lower bandwidth-intensive operations. Such allocation of functions is exemplary only. In one implementation, the high-speed controller **908** is coupled to memory **904**, display **916** (e.g., through a graphics processor or accelerator), and to high-speed expansion ports **910**, which may accept various expansion cards (not shown). In the implementation, low-speed controller **912** is coupled to storage device **906** and low-speed expansion port **914**. The low-speed expansion port, which may include various communication ports (e.g., USB, Bluetooth, Ethernet, wireless Ethernet) may be coupled to one or more input/output devices, such as a keyboard, a pointing device, a scanner, or a networking device such

as a switch or router, e.g., through a network adapter.

(94) The computing device **900** may be implemented in a number of different forms, as shown in the figure. For example, it may be implemented as a standard server **920**, or multiple times in a group of such servers. It may also be implemented as part of a rack server system **924**. In addition, it may be implemented in a personal computer such as a laptop computer **922**. Alternatively, components from computing device **900** may be combined with other components in a mobile device (not shown), such as device **950**. Each of such devices may contain one or more of computing device **900**, **950**, and an entire system may be made up of multiple computing devices **900**, **950** communicating with each other.

(95) Computing device **950** includes a processor **952**, memory **964**, an input/output device such as a display **954**, a communication interface **966**, and a transceiver **968**, among other components. The device **950** may also be provided with a storage device, such as a microdrive or other device, to provide additional storage. Each of the components **950**, **952**, **964**, **954**, **966**, and **968**, are interconnected using various buses, and several of the components may be mounted on a common motherboard or in other manners as appropriate.

(96) The processor **952** may execute instructions within the computing device **950**, including instructions stored in the memory **964**. The processor may be implemented as a chipset of chips that include separate and multiple analog and digital processors. The processor may provide, for example, for coordination of the other components of the device **950**, such as control of user interfaces, applications run by device **950**, and wireless communication by device **950**.

(97) Processor **952** may communicate with a user through control interface **958** and display interface **956** coupled to a display **954**. The display **954** may be, for example, a TFT LCD (Thin-Film-Transistor Liquid Crystal Display) or an OLED (Organic Light Emitting Diode) display, or other appropriate display technology. The display interface **956** may comprise appropriate circuitry for driving the display **954** to present graphical and other information to a user. The control interface **958** may receive commands from a user and convert them for submission to the processor **952**. In addition, an external interface **962** may be provide in communication with processor **952**, so as to enable near area communication of device **950** with other devices. External interface **962** may provide, for example, for wired communication in some implementations, or for wireless communication in other implementations, and multiple interfaces may also be used.

(98) The memory **964** stores information within the computing device **950**. The memory **964** may be implemented as one or more of a computer-readable medium or media, a volatile memory unit or units, or a non-volatile memory unit or units. Expansion memory **974** may also be provided and connected to device **950** through expansion interface **972**, which may include, for instance, a SIMM (Single In Line Memory Module) card interface. Such expansion memory **974** may provide extra storage space for device **950**, or may also store applications or other information for device **950**. Specifically, expansion memory **974** may include instructions to carry out or supplement the processes described above, and may include secure information also. Thus, for example, expansion memory **974** may be provide as a security module for device **950**, and may be programmed with instructions that permit secure use of device **950**. In addition, secure applications may be provided via the SIMM cards, along with additional information, such as placing identifying information on the SIMM card in a non-hackable manner.

(99) The memory may include, for example, flash memory and/or NVRAM memory, as discussed below. In one implementation, a computer program product is tangibly embodied in an information carrier. The computer program product contains instructions that, when executed, perform one or more methods, such as those described above. The information carrier is a computer- or machine-readable medium, such as the memory **964**, expansion memory **974**, memory on processor **952**, or a propagated signal that may be received, for example, over transceiver **968** or external interface **962**.

(100) Device **950** may communicate wirelessly through communication interface **966**, which may

include digital signal processing circuitry where necessary. Communication interface **966** may provide for communications under various modes or protocols, such as GSM voice calls, SMS, EMS, or MMS messaging, CDMA, TDMA, PDC, WCDMA, CDMA2000, or GPRS, among others. Such communication may occur, for example, through radio-frequency transceiver **968**. In addition, short-range communication may occur, such as using a Bluetooth, WiFi, or other such transceiver (not shown). In addition, GPS (Global Positioning System) receiver module **970** may provide additional navigation- and location-related wireless data to device **950**, which may be used as appropriate by applications running on device **950**.

(101) Device **950** may also communicate audibly using audio codec **960**, which may receive spoken information from a user and convert it to usable digital information. Audio codec **960** may likewise generate audible sound for a user, such as through a speaker, e.g., in a handset of device **950**. Such sound may include sound from voice telephone calls, may include recorded sound (e.g., voice messages, music files, etc.) and may also include sound generated by applications operating on device **950**.

(102) The computing device **950** may be implemented in a number of different forms, as shown in the figure. For example, it may be implemented as a cellular telephone **980**. It may also be implemented as part of a smartphone **982**, personal digital assistant, or other similar mobile device.

(103) Device **950** may also include one or more different devices that are capable of sensing motion. Examples include, but are not limited to, accelerometers and compasses. Accelerometers and compasses, or other devices that are capable of detecting motion or position are available from any number of vendors and may sense motion in a variety of ways. For example, accelerometers may detect changes in acceleration while compasses may detect changes in orientation relative to the magnetic North or South Pole. These changes in motion may be detected by the device **950** and used to update the display of the respective devices **950** according to processes and techniques described herein.

(104) Various implementations of the systems and techniques described here may be realized in digital electronic circuitry, integrated circuitry, specially designed ASICs (application specific integrated circuits), computer hardware, firmware, software, and/or combinations thereof. These various implementations may include implementation in one or more computer programs that are executable and/or interpretable on a programmable system including at least one programmable processor, which may be special or general purpose, coupled to receive data and instructions from, and to transmit data and instructions to, a storage system, at least one input device, and at least one output device.

(105) These computer programs (also known as programs, software, software applications or code) include machine instructions for a programmable processor, and may be implemented in a high-level procedural and/or object-oriented programming language, and/or in assembly/machine language. As used herein, the terms “machine-readable medium” “computer-readable medium” refers to any computer program product, apparatus and/or device (e.g., magnetic discs, optical disks, memory, Programmable Logic Devices (PLDs)) used to provide machine instructions and/or data to a programmable processor, including a machine-readable medium that receives machine instructions as a machine-readable signal. The term “machine-readable signal” refers to any signal used to provide machine instructions and/or data to a programmable processor.

(106) To provide for interaction with a user, the systems and techniques described here may be implemented on a computer having a display device (e.g., a CRT (cathode ray tube) or LCD (liquid crystal display) monitor) for displaying information to the user and a keyboard and a pointing device (e.g., a mouse or a trackball) by which the user may provide input to the computer. Other kinds of devices may be used to provide for interaction with a user as well; for example, feedback provided to the user may be any form of sensory feedback (e.g., visual feedback, auditory feedback, or tactile feedback); and input from the user may be received in any form, including acoustic, speech, or tactile input.

(107) The systems and techniques described here may be implemented in a computing system that includes a back end component (e.g., as a data server), or that includes a middleware component (e.g., an application server), or that includes a front end component (e.g., a client computer having a graphical user interface or a Web browser through which a user may interact with an implementation of the systems and techniques described here), or any combination of such back end, middleware, or front end components. The components of the system may be interconnected by any form or medium of digital data communication (e.g., a communication network). Examples of communication networks include a local area network ("LAN"), a wide area network ("WAN"), and the Internet.

(108) The computing system may include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other.

(109) FIG. **10** is a block diagram of an example data structure **1000** of a language model. The data structure **1000** can weightedly link a semantic idea or category with language models, for example, for use in speech recognition.

(110) A semantic category **1002**, such as a search query or type of input field, can be linked to one or more of a collection of interpolated language models **1004**. The interpolated language models **1004** can be linked to one or more base language models **1006**. The links between the interpolated language models **1004** and base language models **1006** can be weighted. In some examples, the sum of the weights of the links from one interpolated language model can be 1 or can be normalized to 1.

(111) It will be understood that, although a particular number and configuration of interpolated language models **1004**, base language models **1006**, and links are shown, other numbers and configurations are possible. For example, sufficient interpolated language models **1004** may exist that every weighted combination of base language models **1006** has a linked interpolated language model **1004**. In some examples, interpolated language model can be linked to more or fewer base language models **1006**. In some examples, different link weights between interpolated language models **1004** and base language models **1006** may exist, such as positive integers, probabilities, or dimensional distance (e.g. W, X, Y, Z values for four dimensional space.) In some implementations, multiple semantic categories **1002** can be linked to a single interpolated language model **1004**.

(112) In some examples, interpolated language models **1004** can be created on demand, such as when a semantic category **1002** is added to the data structure. The interpolated language models **1004** can persists after the removal of semantic categories (not shown). For example, previously removed semantic categories (not shown) may have prompted the creation of the interpolated language models **1004a**, **1004b**, and **1004c**. A new semantical category **1008** can be added to the data structure **1004**, and linked to any of the interpolated language models **1004**.

(113) In some examples, every possible interpolated language model **1004** can be pre-created for use by a new semantic category **1008**. Some of these examples can be used in situations, such as when many new semantic categories **1008** are expected, when pre-processing time is available, and/or when few base language models **1006** are expected.

(114) A number of embodiments have been described. Nevertheless, it will be understood that various modifications may be made without departing from the spirit and scope of what is described here.

(115) In addition, the logic flows depicted in the figures do not require the particular order shown, or sequential order, to achieve desirable results. In addition, other steps may be provided, or steps may be eliminated, from the described flows, and other components may be added to, or removed from, the described systems. Accordingly, other embodiments are within the scope of the following claims.

Claims

1. A computer-implemented method executed on data processing hardware of a mobile computing device that causes the data processing hardware to perform operations comprising: displaying, on an electronic display of the mobile computing device, a user interface for an application executing on the mobile computing device, the user interface comprising a text field and a graphical button, the graphical button displayed next to the text field on the electronic display; in response to receiving first interaction data indicating user interaction with the text field, invoking the user interface to further display a virtual keyboard on the electronic display; in response to receiving second interaction data indicating user interaction with the graphical button: removing the virtual keyboard from display on the electronic display so as to disable an ability of the application executing on the mobile computing device to receive typed text input; and enabling a voice input mode for the application executing on the mobile computing device by invoking the user interface to display a visual indication that the application is enabled to receive voice input; and after enabling the voice input mode for the application, detecting a spoken utterance directed toward the application.
2. The computer-implemented method of claim 1, wherein the operations further comprise, after detecting the spoken utterance directed toward the application: providing audio data corresponding to the spoken utterance to a server system, the server system comprising a speech recognition system; receiving, from the server system, a transcription of the spoken utterance generated by the server system; and displaying the transcription of the spoken utterance on the electronic display.
3. The computer-implemented method of claim 2, wherein the operations further comprise: determining a context associated with the spoken utterance; providing, to the server system, data that indicates the context associated with the spoken utterance; and receiving the transcription of the spoken utterance from the server system, wherein the server system generated the transcription based on the data that indicates the context associated with the utterance.
4. The computer-implemented method of claim 1, wherein the spoken utterance is captured by a microphone of the mobile computing device.
5. The computer-implemented method of claim 1, wherein the mobile computing device comprises an audio output device.
6. The computer-implemented method of claim 1, wherein the operations further comprise, in response to detecting the spoken utterance, invoking the user interface to display a graphic indicating that speech-to-text conversion on the spoken utterance is in progress.
7. The computer-implemented method of claim 6, wherein the operations further comprise, in response to detecting the spoken utterance, invoking the user interface to further display a graphical cancel button for canceling the speech-to-text conversion on the spoken utterance.
8. The computer-implemented method of claim 1, wherein the user interface comprises a user interface for a multi-modal input method editor that enables the application executing on the mobile computing device to receive voice input and typed input.
9. The computer-implemented method of claim 1, wherein the operations further comprise, in response to receiving the second interaction data indicating user interaction with the graphical button, removing the graphical microphone from display on the electronic display.
10. The computer-implemented method of claim 1, wherein the mobile computing device comprises a mobile phone.
11. A mobile computing device comprising: an electronic display; data processing hardware; and memory hardware in communication with the data processing hardware and storing instructions that when executed on the data processing hardware cause the data processing hardware to perform operations comprising: displaying, on the electronic display, a user interface for an application executing on the mobile computing device, the user interface comprising a text field and a

graphical button, the graphical button displayed next to the text field on the electronic display; in response to receiving first interaction data indicating user interaction with the text field, invoking the user interface to further display a virtual keyboard on the electronic display; in response to receiving second interaction data indicating user interaction with the graphical button: removing the virtual keyboard from display on the electronic display so as to disable an ability of the application executing on the mobile computing device to receive typed text input; and enabling a voice input mode for the application executing on the mobile computing device by invoking the user interface to display a visual indication that the particular application is enabled to receive voice input; and after enabling the voice input mode for the application, detecting a spoken utterance directed toward the application.

12. The mobile computing device of claim 11, wherein the operations further comprise, after detecting the spoken utterance directed toward the application: providing audio data corresponding to the spoken utterance to a server system, the server system comprising a speech recognition system; receiving, from the server system, a transcription of the spoken utterance generated by the server system; and displaying the transcription of the spoken utterance on the electronic display.

13. The mobile computing device of claim 12, wherein the operations further comprise: determining a context associated with the spoken utterance; providing, to the server system, data that indicates the context associated with the spoken utterance; and receiving the transcription of the spoken utterance from the server system, wherein the server system generated the transcription based on the data that indicates the context associated with the utterance.

14. The mobile computing device of claim 11, wherein the spoken utterance is captured by a microphone of the mobile computing device.

15. The mobile computing device of claim 11, wherein the mobile computing device comprises an audio output device.

16. The mobile computing device of claim 11, wherein the operations further comprise, in response to detecting the spoken utterance, invoking the user interface to display a graphic indicating that speech-to-text conversion on the spoken utterance is in progress.

17. The mobile computing device of claim 16, wherein the operations further comprise, in response to detecting the spoken utterance, invoking the user interface to further display a graphical cancel button for canceling the speech-to-text conversion on the spoken utterance.

18. The mobile computing device of claim 11, wherein the user interface comprises a user interface for a multi-modal input method editor that enables the application executing on the mobile computing device to receive voice input and typed input.

19. The mobile computing device of claim 11, wherein the operations further comprise, in response to receiving the second interaction data indicating user interaction with the graphical button, removing the graphical microphone from display on the electronic display.

20. The mobile computing device of claim 11, wherein the mobile computing device comprises a mobile phone.
