

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386586
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Doan; Nhat Vu et al.

Speech processing and multi-modal widgets

Abstract

Techniques for performing speech processing using multi-modal widget information are described. A system may receive input data corresponding to a user input. The system may also receive widget context data corresponding to one or more multi-modal widgets active at a device. The system may use the widget context data to perform natural language understanding (NLU) processing with respect to the user input, and for selecting a skill component for responding to the user input. The system may send a widget identifier to the skill component when invoking the skill to respond to the user input.

Inventors: Doan; Nhat Vu (Seattle, WA), Cummings; Nicholas Adam (Lynnwood, WA), Thakare; Prashant Jayaram (Sammamish, WA), Kumar; Jalaj (Seattle, WA), Ravi; Ganesh Prabu (Seattle, WA), Wang; Chih-Shin (Bellevue, WA), Gyanchandani; Narenda (Sammamish, WA)

Applicant: Amazon Technologies, Inc. (Seattle, WA)

Family ID: 91227773

Assignee: Amazon Technologies, Inc. (Seattle, WA)

Appl. No.: 18/631683

Filed: April 10, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240256219 A1	Aug. 01, 2024

Related U.S. Application Data

continuation parent-doc US 17488385 20210929 US 11966663 child-doc US 18631683

Publication Classification

Int. Cl.: **G10L15/00** (20130101); **G06F3/16** (20060101); **G10L15/197** (20130101); **G10L15/22** (20060101)

U.S. Cl.:

CPC **G06F3/167** (20130101); **G10L15/197** (20130101); **G10L15/22** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9262120	12/2015	Chipchase	N/A	G06F 3/167
10395655	12/2018	Kumar	N/A	G06N 5/022
11094316	12/2020	Visser	N/A	G10L 17/04
11966663	12/2023	Doan	N/A	G10L 15/1815
2013/0144629	12/2012	Johnston	704/E11.001	G10L 15/22
2020/0058300	12/2019	Kumar	N/A	G10L 15/1815
2020/0257734	12/2019	Steinfort	N/A	G06F 9/451
2021/0012770	12/2020	Choudhary	N/A	G06F 3/017
2021/0182016	12/2020	Ding	N/A	H04R 29/001
2021/0182018	12/2020	Szmigiel	N/A	G06F 3/167

Primary Examiner: Singh; Satwant K

Attorney, Agent or Firm: Pierce Atwood LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION DATA (1) This application is a continuation of, and claims priority to, U.S. patent application Ser. No. 17/488,385, entitled “SPEECH PROCESSING AND MULTI-MODAL WIDGETS,” filed on Sep. 29, 2021, the contents which are expressly incorporated herein by reference in their entirety.

BACKGROUND

(1) Natural language processing systems have progressed to the point where humans can interact with and control computing devices using their voices. Such systems employ techniques to identify the words spoken by a user based on the various qualities of received input data. Speech recognition combined with natural language understanding processing techniques enable speech-based user control of computing devices to perform tasks based on the spoken inputs. Speech recognition and natural language understanding processing techniques are sometimes referred to collectively or separately as spoken language understanding (SLU) processing. SLU processing

may be used by computers, hand-held devices, telephone computer systems, kiosks, and a wide variety of other devices to improve human-computer interactions.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) For a more complete understanding of the present disclosure, reference is now made to the following description taken in conjunction with the accompanying drawings.
- (2) FIG. 1 is a conceptual diagram illustrating a system for performing speech processing in view of multi-modal widgets presented at a device, according to embodiments of the present disclosure.
- (3) FIG. 2 conceptually illustrates example widget context data, according to embodiments of the present disclosure.
- (4) FIG. 3 is a conceptual diagram illustrating components that may be included in a device, according to embodiments of the present disclosure.
- (5) FIG. 4 is a conceptual diagram of how natural language understanding (NLU) processing may be performed, according to embodiments of the present disclosure.
- (6) FIG. 5 is a conceptual diagram illustrating a skill selection component configured to select a skill for responding to a user input, according to embodiments of the present disclosure.
- (7) FIG. 6 is a block diagram conceptually illustrating example components of a device, according to embodiments of the present disclosure.
- (8) FIG. 7 is a block diagram conceptually illustrating example components of a system, according to embodiments of the present disclosure.
- (9) FIG. 8 illustrates an example of a computer network for use with the overall system, according to embodiments of the present disclosure.

DETAILED DESCRIPTION

- (10) Automatic speech recognition (ASR) is a field of computer science, artificial intelligence, and linguistics concerned with transforming audio data associated with speech into text representative of that speech. Similarly, natural language understanding (NLU) is a field of computer science, artificial intelligence, and linguistics concerned with enabling computers to derive meaning from text input containing natural language. ASR and NLU are often used together as part of a speech processing system. Text-to-speech (TTS) is a field of concerning transforming textual data into audio data that is synthesized to resemble human speech.
- (11) Certain devices that have a display screen may present content to a user. Some such devices may be configured to present content from multiple applications via, what is referred to as, multi-modal widgets herein. A multi-modal widget may correspond to a portion of a display screen that presents visual content, that may automatically updates when appropriate, and that the user may interact with using touch inputs and/or voice inputs.
- (12) Multi-modal widgets, as used herein, may be glanceable, self-updating, and interactive view of content and functionality displayed on a device screen. Multi-modal widgets may offer ambient experiences (where various technology, systems, devices—and data gathered by them—seamlessly interact and adapt to a user's needs), enable better multitasking on the system, and increase the discoverability of skills. For example, a developer can use the multi-modal widgets to surface latest content or upcoming events, interact with users without starting a full skill experience, and give users a quick way to go directly to the desired or most interesting part of a skill. In addition to viewing content, a user may also perform actions using a multi-modal widget. For example, a user can check an item off a list or add an item to a list using touch or voice inputs. Depending on the user action, the multi-modal widget may update inline (i.e., update the present content at the multi-modal widget) or start a full-screen experience (i.e., open a new user interface screen).
- (13) A developer may create a multi-modal widget by defining a rendered document (via software

code) and how information is to be presented visually via the rendered document, how instructions are to be executed with respect to the rendered document, etc. The developer may also define data sources from which the multi-modal widget can be populated. For example, a weather multi-modal widget may receive weather data from a weather service. A single multi-modal widget may receive data from multiple sources. The multi-modal widget may update the data by requesting updated data from the data source or based on the data source pushing updated data to the multi-modal widget. The rendered document definition may include instructions on how the data is to be updated.

(14) A system may control a multi-modal widget at a device using a particular type of software language. In some embodiments, the system may control multi-modal widgets using a JSON-based HTML5 language. An example of such a software language is the Alexa Presentation Language (APL). Other types of software languages may be used to control multi-modal widgets. Using APL, a developer may create visual experiences to accompany a skill, such as, animations, graphics, images, slideshows, video, etc., which may be presented to a user via a multi-modal widget. A developer may create JSON files including software instructions to control the multi-modal widget. Such JSON files may be referred to herein as an APL document. The APL document may be invoked and downloaded to a user device. The user device may import images and other data indicated in the APL document and render the programmed experience at the user device.

(15) The system may provide multi-modal widget templates that can be customized by a developer using the APL document. The multi-modal widget templates may correspond to different widget sizes that can be used to control the size of the multi-modal widget on the device screen. The multi-modal widgets templates may include a list template that can be used to display a list of items (e.g., a shopping list, a task list, etc.), which may include text and/or images. The multi-modal widgets templates may include a text wrapping template that can be used for text-based experiences (e.g., displaying tips, facts, instructions, etc.). Another multi-modal template may be an action button template that can be used to present content along with a button that a user can select to perform an indicated action. Another multi-modal template may be an image-caption template that can be used to present an image along with text. Yet another multi-modal template may be a photo template that can be used to present an image as focused-content.

(16) Some devices may enable a user to interact with skills. As used herein, a “skill” may refer to software, that may be installed on a machine or a virtual machine (e.g., software that may be launched in a virtual instance when called or otherwise activated), configured to process natural language understanding (NLU) output data (e.g., including an intent and optionally one or more entities) and perform one or more actions in response thereto. What is referred to herein as a skill may sometimes be referred to as or otherwise incorporated into an application, bot, action, or the like. A group of skills of related functionality may be associated with a domain. For example, a first music skill and a second music skill may be associated with a music domain. A user may interact with skills using voice inputs and/or touch inputs. Skills may provide multi-modal widgets to enable the user to interact with the skill using touch inputs.

(17) The present disclosure relates to performing speech processing in view of multi-modal widgets presented by a screen of a device. Techniques of the present disclosure enable a user to multitask and seamlessly switch between interactions with multiple multi-modal widgets. Techniques of the present disclosure also enables a user to switch between interacting with a skill corresponding to a multi-modal widget and another skill that may not have a corresponding multi-modal widget presented at the device.

(18) A device or a system, of the present disclosure, can use contextual information of the multi-modal widgets being presented at the device to interpret natural language inputs provided by the user. The widget context data may indicate the content presently being displayed at the multi-modal widget. The widget context data may also include a name of the multi-modal widget, a skill associated with the multi-modal widget, a widget identifier and a widget instance identifier. The

device/system may use the widget context data to determine an intent of a natural language input and entity data corresponding to the natural language input. For example, the device/system may determine, using the widget context data, that the intent of the natural language input corresponds to the multi-modal widget or that the entity data corresponds to the content being presently displayed at the multi-modal widget. The device/system may further use the widget context data to select a skill for processing with respect to the natural language input.

(19) Multiple instances, presenting different content, of the same multi-modal widget may be presented at a device. For example, a device may display weather information for a first city via a first instance of a weather multi-modal widget, and the device may also display weather information for a second city via a second instance of the weather multi-modal widget. In some embodiments, the device/system determines which multi-modal widget instance, at the device, the natural language input corresponds to, and send an indication of the multi-modal widget instance to the skill. The skill may use the multi-modal widget instance to determine an output responsive to the natural language input. The skill may alternatively or additionally send a command to the device to present an output responsive to the natural language input via the multi-modal widget instance. For example, the device may present a first shopping list via a first instance of a list multi-modal widget and a second shopping list via a second instance of the list multi-modal widget. The device/system may determine that a user's voice input to: "remove [item] from my list" corresponds to the first shopping list, and may send an identifier for the first instance of the list multi-modal widget to a list skill, so that the list skill may perform an action, responsive to the user input, with respect to the first list. As part of the action, the list skill may cause the device to update the first instance of the list multi-modal widget (corresponding to the first shopping list) based on the received identifier for the first instance of the list multi-modal widget.

(20) Using the widget context data, skills may determine how an output is to be presented to the user. For example, a skill may select from presenting an output within the widget's existing layout/display configuration (referred to herein as "inline") or launching a full screen interface to present the output (i.e. opening a new screen/window to present the output). The skill may determine how to present the output based on the widget context data indicating the content being presently displayed at the device at the multi-modal widget. For example, if the output relates to the presently displayed content, then the skill may present the output inline. As another example, if the output does not relate to the presently displayed content, then the skill may present the output via a full-screen interface.

(21) Teachings of the present disclosure may be configured to incorporate user permissions and may only be performed if approved by a user. As such, the systems, devices, components, and techniques described herein would be typically configured to restrict processing where appropriate and only process user data in a manner that ensures compliance with all appropriate laws, regulations, standards, and the like. The teachings of the present disclosure can be implemented on a geographic basis to ensure compliance with laws in various jurisdictions and entities in which the computing components and/or user are located.

(22) FIG. 1 shows a system **100** configured to perform speech processing in view of multi-modal widgets presented at a device. As shown in FIG. 1, the system **100** may include a device **110**, local to a user **105**, and a system **120** connected across one or more networks **199**. The network(s) **199** may include the Internet and/or any other wide- or local-area network, and may include wired, wireless, and/or cellular network hardware. Although the figures and discussion of the present disclosure illustrate certain steps in a particular order, the steps described may be performed in a different order (as well as certain steps removed or added) without departing from the present disclosure.

(23) The system **120** may be a speech processing system configured to process spoken natural language inputs using ASR and NLU processing. The system **120** may include multiple components to facilitate speech processing, such as, an orchestrator component **130**, an ASR

component **150**, a NLU component **160**, a skill selection component **165**, a skill request router **188**, and one or more skill components **190**. The system **120** may also include a profile storage **170**, a TTS component **180**, and a user recognition component **195** to facilitate processing of users inputs and generating outputs. One or more of the skill components **190** may be in (wired or wireless) communication with a skill system(s) **125** located remote/external to the system(s) **120**.

(24) The device **110** may display content via one or more multi-modal widgets **112**. The content may include text, icons, images, video, selectable user interface elements (e.g., buttons, links, etc.), and the like. Each widget **112** may be associated with a same or different skill component **190**. Content at the widget **112** may be determined (including updated, refreshed, etc.) by the corresponding skill component **190**.

(25) The user **105** may select which multi-modal widgets are to display content at the device **110**. The user **105** may select certain multi-modal widgets **112** to be active and presenting content continuously (e.g., on a daily basis). For example, the user **105** may select a weather multi-modal widget to be active daily, and the weather multi-modal widget may update weather content presented via the weather multi-modal widget. Such multi-modal widgets **112** may automatically activate after a reboot of the device **110**.

(26) Referring to FIG. 1, the user **105** may speak an input, and the device **110** may capture audio **107** representing the spoken input. For example, the user **105** may say “Alexa, add [item] to my shopping list” or “Alexa, show me weather for [city].” In other examples, the user **105** may provide another type of input (e.g., selection of a button, selection of displayed graphical interface elements, may perform a gesture, etc.). The device **110** may send (step 1) audio data (or other type of input data, such as, image data, text data, etc.) corresponding to the user input to the system **120** for processing. In particular, the orchestrator component **130** may receive the input data from the device **110**.

(27) In response to receiving the user input from the user **105**, the device **110** may also send (step 2) widget context data **114** to the system **120**. The widget context data **114** may correspond to the multi-modal widgets **112** active at the device **110** (i.e., presenting content via the device **110** when the device **110** receives the user input). The widget context data **114** may include widget identifiers for the multi-modal widgets **112**. In some embodiments, the widget context data **114** may also include a skill identifier associated with the skill component **190** corresponding to the multi-modal widget **112**. For example, the widget context data **114** may include a first widget identifier for the multi-modal widget **112a**, a first skill identifier corresponding to the multi-modal widget **112a**, a second widget identifier for the multi-modal widget **112b**, and a second skill identifier corresponding to the multi-modal widget **112b**. FIG. 2 illustrates example data included in the widget context data **114** corresponding to the multi-modal widget **112**. In some embodiments, the device **110** may send first widget context data **114a** corresponding to the first multi-modal widget **112a**, and may send separate second widget context data **114b** corresponding to the second multi-modal widget **112b**.

(28) In some embodiments, the widget context data **114** may also include a time when the user **105** most recently/last interacted with the multi-modal widget **112** (as illustrated in FIG. 2). This time may be represented, in some embodiments, as time elapsed since the last interaction. The device **110** may store such information based on receiving user inputs (e.g., touch inputs, voice inputs, etc.) from the user **105** corresponding to the multi-modal widget **112**.

(29) In some embodiments, the widget context data **114** may also include a time when the multi-modal widget **112** was last updated (as illustrated in FIG. 2). This time may be based on the last time when content was updated at the multi-modal widget **112** or other data corresponding to the multi-modal widget **112** was updated (e.g., update multi-modal widget software, update multi-modal widget interface, etc.). For example, this time may be based on a last time when the device **110** received a command, from a skill component **190**, to update the multi-modal widget **112**.

(30) In some embodiments, the widget context data **114** may also include information on where the

multi-modal widget **112** is positioned (referred to as position information herein) on the display screen of the device **110** (as illustrated in FIG. 2). The position information may be indicated in terms of quadrants or portions. For example, the position information for the multi-modal widget **112a** (in FIG. 1) may be “left side panel”, and the position information for the multi-modal widget **112b** (in FIG. 1) may be “center” or “bottom half” Other position information may be right side panel, top panel, bottom panel, center, bottom half, top half, etc. In other embodiments, the position information may be indicated as coordinates along x-axis and y-axis (e.g., x-y coordinates) or pixel coordinates corresponding to the display interface of the device **110**. The coordinates may be a set of four coordinates, each one corresponding to a corner of the multi-modal widget.

(31) The content displayed via the multi-modal widget **112** may be scrollable (i.e. the user **105** may be able to scroll up, scroll down, scroll left, or scroll right) to view different portions of the content. In some embodiments, the widget context data **114** may also include scroll information (as illustrated in FIG. 2).

(32) The widget context data **114**, in some embodiments, may also include content (as illustrated in FIG. 2) being displayed via the multi-modal widget **112** at the device **110** at the time the user input from the user **105** is received. Some multi-modal widgets **112** may not display an entirety of the content at the device **110** due to multi-modal widget size, screen size, etc. As such, the multi-modal widget **112** may display a portion of the content based on the multi-modal widget size and the screen size. For example, some multi-modal widgets may have multiple pages or tabs that the user **105** can select to view other portions of the content. As a further example, some multi-modal widgets may have a scroll bar so the user **105** can view other portions of the content. The widget context data **114** may include (or otherwise indicate) the portion of the content being displayed at the device **110**, as it is viewable by the user **105**.

(33) The widget context data **114** may, in some embodiments, also include a layout type of the multi-modal widget **112** at the time the user input is received (as illustrated in FIG. 2). The layout type may be full screen (when the multi-modal widget is in a full-screen mode), overlay (when the multi-modal widget is overlaid over other multi-modal widgets), non-full screen, etc.

(34) The widget context data **114**, in some embodiments, may also include a focus type for the multi-modal widget **112** (as illustrated in FIG. 2). The focus type may be touch focus (when the user **105** touches the multi-modal widget **112** before or while providing the user input), visual focus (when the multi-modal widget **112** is visible on the display screen of the device **110** while the user **105** provides the user input), or non-visual focus (when the multi-modal widget **112** is active but not visible on the display screen of the device **110**, and may be displayed on a different tab or page of the display screen of the device **110**, or may be at least partially obscured by an overlaid multi-modal widget).

(35) The widget context data **114**, in some embodiments, may also include a user-defined name for the multi-modal widget **112** (as illustrated in FIG. 2). For example, the user **105** may name a shopping list multi-modal widget as [list 1]. The user **105** may provide a spoken input or a touch input to name the multi-modal widget **112**.

(36) In some embodiments, the information described above as being included in the widget context data **114** may be provided by the device **110**. In other embodiments, some or all of the above described information may be determined by the system **120**. For example, the device **110** may send the widget identifiers for the multi-modal widgets **112** presented at the device **110**, and the system **120** may determine the content being presented via the multi-modal widgets **112**. The system **120** may send a request to the skill component **190** corresponding to the multi-modal widget **112** to determine the content being presented via the multi-modal widget **112** associated with the widget identifier.

(37) The orchestrator component **130** may store (step 3) the widget context data **114** in a widget session storage **142**. The widget context data **114**, in the widget session storage **142**, may be associated with a device identifier for the device **110** and/or a user identifier for the user **105**. The

user identifier for the user **105** may be determined by the user recognition component **195** as described in detail below. The device identifier may be sent by the device **110** along with the widget context data **114**, or with the input data corresponding to the user input from the user **105**. (38) The widget session storage **142** may store historic widget context data corresponding to the user identifier associated with the user **105**, and for user identifiers associated with other users. The historic widget context data may correspond to one or more multi-modal widgets that may have been previously active at the device **110**, and may not be active currently. The historic widget context data may correspond to one or more multi-modal widgets that may have been enabled for the device **110**, and may be currently disabled. In addition to the widget context data, the widget session storage **142** may also store interaction data corresponding to a multi-modal widget. Such interaction data may relate to user inputs received at the system **120** and corresponding to the multi-modal widget, output data generated by a skill component and corresponding to the multi-modal widget (i.e. output data presented via the multi-modal widget), NLU data corresponding to the user inputs, selected skill identifier, time the user input was received, etc. The widget session storage **142** may associate the widget context data and the corresponding interaction data with a widget session identifier in the storage **142**. One or more components of the system **120** may use the data stored at the widget session storage **142** to perform its processing. For example, the skill selection component **165** may use the data stored at the widget session storage **142** (and corresponding to historic interactions between the user **105** and multi-modal widgets) to select a skill component for the instant user input (received in step **1**). As another example, the NLU component **160** may use the data stored at the widget session storage **142** to determine NLU data corresponding to the instant user input.

(39) In the case that the input data (received in step **1**) is audio data, the orchestrator component **130** may send (step **4**) the audio data to the ASR component **150**, and the ASR component **150** may process the audio data to determine ASR data (e.g., token data, text data, one or more ASR hypotheses including token or text data and corresponding confidence scores, etc.) corresponding to the words spoken by the user **105**. Details on how the ASR component **150** may process the audio data are described below. The ASR component may send (step **5**) the ASR data to the orchestrator component **130**.

(40) In some embodiments, the orchestrator component **130** may also send (at step **4**) the widget context data **114** to the ASR component **150**. The ASR component **150** may use the widget context data **114** to perform ASR processing. For example, the ASR component **150** may employ one technique involving boosting of one or more words (represented as text or token data) included in the widget context data **114**, such as, the name of the multi-modal widget **112**, words included in the content presented at the multi-modal widget **112**, etc. Using such as a technique may enable the ASR component **150** to recognize certain rare or personalized words that may be included in the spoken input from the user **105**. In some embodiments, the ASR component **150** may determine embedding data corresponding to the words included in the widget context data **114**, and use the embedding data to determine the ASR data corresponding to the audio data.

(41) The orchestrator component **130** may send (step **6**) the ASR data and the widget context data **114** to the NLU component **160**. The NLU component **160** may determine NLU data corresponding to the user input, where the NLU data may include an intent and one or more entities. The NLU component **160** may use the widget context data **114** to determine the intent and/or the one or more entities. In some cases, the NLU component **160** may determine the intent based on which multi-modal widgets **112** are presented at the device **110**. In some cases, the NLU component **160** may determine an entity based on the content being displayed at the multi-modal widget **112**. In some cases, the NLU component **160** may determine the NLU data (intent and/or entity) based on the focus information for the multi-modal widget **112**, content being displayed via the multi-modal widget **112**, time of last interaction with the multi-modal widget **112**, etc. For example, the device **110** may present a music multi-modal widget displaying information about a [song name], and the

user **105** may say “Play song.” In this example, the NLU component **160** may determine, using the ASR data corresponding to the spoken input and the widget context data **114**, the following NLU data: {intent: <PlayMusic>; entity type: <SongName> entity: “[song name]”}. The entity [song name] may be based on the widget context data **114** indicating that the multi-modal widget **112** displaying information about the [song name]. In some embodiments, the widget context data **114** may be used to target certain Finite-state Transducers (FSTs) implemented by the NLU component **160**. For example, the NLU component **160** may implement FSTs specific to certain skills, and based on the widget context data **114** indicating which multi-modal widgets **112**, and in turn which skills, are active/enabled at the device **110**, the NLU component **160** may use or boost the corresponding FSTs. Further details about how the NLU component **160** determines NLU data are described below in relation to FIG. 4. The NLU component **160** may send (step 7) the NLU data, corresponding to the user input, to the orchestrator component **130**.

(42) The orchestrator component **130** may send (step 8) the NLU data and the widget context data **114** to the skill selection component **165**. The skill selection component **165** may be configured to determine which skill component **190** is capable of responding to the user input. The skill selection component **165** may make this determination based on which skill component **190** is capable of processing the intent and the entity data included in the NLU data. The skill selection component **165** may make this determination further based on the information included in the widget context data **114**. In selecting the skill component **190**, the skill selection component **165** may determine which particular multi-modal widget **112** the user input corresponds to. Based on determining which multi-modal widget **112** the user input corresponds to, the skill selection component **165** may select the skill component **190** to respond to the user input. In some embodiments, the skill selection component **165** may select the skill component **190** associated with the skill identifier included in the widget context data **114** as corresponding to the multi-modal widget **112**. For example, if the skill selection component **165** determines that the user input corresponds to the multi-modal widget **112a**, then the skill selection component **165** may select the skill component **190a** associated with the first skill identifier corresponding to the multi-modal widget **112a** (as indicated in the widget context data **114**).

(43) The skill selection component **165** may determine which multi-modal widget **112** corresponds to the user input based on the intent and/or the entity included in the NLU data. For example, if the intent is <PlayMusic>, then the skill selection component **165** may determine that the music multi-modal widget, presented at the device **110**, corresponds to the user input.

(44) In some cases, the device **110** may present multiple instances of the same multi-modal widget. In some embodiments, the skill selection component **165** may determine which multi-modal widget instance corresponds to the user input. The skill selection component **165** may make this determination based on the content displayed via each multi-modal widget instance, the content that is visible on the device screen when the user **105** provided the user input, the time since last interaction, the focus information for the multi-modal widget, user-defined name, and/or other information included in the widget context data **114**. For example, the multi-modal widgets **112a** and **112b** may be different instances of a shopping list multi-modal widget, and the user **105** may say “remove [item] from my list.” In this example, the skill selection component **165** may determine that the content displayed via the multi-modal widget **112a** includes the [item], and may select the multi-modal widget **112a** as corresponding to the user input. In another example, a first list at a first instance of a shopping list multi-modal widget may be named by the user **105** [list 1], a second list at a second instance of the shopping list multi-modal widget may be named by the user **105** [list 2], and the user **105** may say “add [item] to my [list 1].” In this example, the skill selection component **165** may determine that the user input corresponds to the first shopping list multi-modal widget.

(45) In some embodiments, the skill selection component **165** may output a widget identifier associated with the multi-modal widget **112** that is determined to correspond to the user input. In

some embodiments, the skill selection component **165** may output more than one widget identifier if the skill selection component **165** determines the user input corresponds to more than one multi-modal widget **112**.

(46) The skill selection component **165** may send (step **9**) a skill identifier associated with the skill component **190** to the orchestrator component **130**. In some embodiments, the skill selection component **165** may also send a widget identifier associated with the multi-modal widget **112** determined to correspond to the user input.

(47) The orchestrator component **130** may send (step **10**) the skill identifier, the widget identifier and the NLU data to the skill request router component **188**. The orchestrator component **130** may also send a command to the skill request router component **188** to send the NLU data to the skill component **190** associated with the skill identifier, so that the skill component **190** may process with respect to the user input.

(48) In some embodiments, the skill request router component **188** may determine whether data, to be sent to the skill component **190** and indicated by the orchestrator component **130**, is appropriate to send to the skill component **190**. As described herein, the system **120** may implement more than one skill component **190**, each of which may be configured to perform certain functionalities. When a particular skill component **190a** is to be invoked for a user input, the other skill components **190** are not to be invoked and are not to receive the NLU data or other data corresponding to the user input.

(49) In some embodiments, the orchestrator component **130** may also send the widget context data **114** to the skill request router component **188** for sending to the skill component **190**. The skill request router component **188** may determine which portion of the widget context data **114** is to be sent to the skill component **190**. As described herein, the widget context data **114** may include data corresponding to all multi-modal widgets **112** active or enabled at the device **110**. The skill request router component **188** may determine widget context data **116** to be a portion of the widget context data **114** that corresponds to the widget identifier and/or the skill identifier received from the orchestrator component **130** at step **10**. In this manner, the skill request router **188** may not send, to the skill component **190**, all the content being presented at the device **110**. This may prevent the skill component **190** from being able to identify content being presented at the device **110** and which the skill component **190** does not need to process (e.g., multi-modal widget content corresponding to a different skill component).

(50) As used herein, a multi-modal widget may be “active” when the multi-modal widget is presenting content at the device **110**. In other words, when the user **105** adds the multi-modal widget to a GUI (e.g., home screen, another tab, etc.) of the device **110**. In some cases, the GUI of the device **110** may include multiple tabs or pages that the user **105** can scroll or switch to view other content. The user **105** may add one or more multi-modal widgets to a first tab/page, and add other multi-modal widgets to another tab/page.

(51) As used herein, a multi-modal widget may be enabled/installed when the user **105** downloads a multi-modal widget to the device **110**. Such enabled/installed multi-modal widgets may not present content via the device **110** until the user **105** adds the multi-modal widget to the GUI of the device **110**.

(52) In some embodiments, the functionalities described herein may take into consideration active multi-modal widgets to determine NLU data corresponding to a user input and select a skill component **190** to respond to the user input. In some embodiments, the functionalities described herein may also take into consideration enabled/installed multi-modal widgets. For example, the user **105** may say “add [multi-modal widget name] to the home screen,” in which case the NLU component **160** may determine entity data to be [multi-modal widget name] which may correspond to one of the enabled/installed multi-modal widgets, and the skill selection component **165** may select the skill component **190** corresponding to the particular enabled/installed multi-modal widget.

(53) In some embodiments, the skill component **190** may store/track content data being presented via the multi-modal widget **112** at the device **110**. For example, the skill component **190** may store the content data when the content data is sent (via the orchestrator **130**) to the device **110** for output via the multi-modal widget **112**. In some cases, the multi-modal widget **112** may be scrollable or may update dynamically (e.g., update content based on time, etc.). In such cases, the content presently displayed at the device **110** and presently visible to the user **105** may be a portion of the content data that the skill component **190** is storing/tracking. For example, the skill component **190** may send weather data corresponding to multiple days for display via a weather multi-modal widget. The weather multi-modal widget may have a first view/tab that displays weather for the present day, a second view/tab that displays weather for the next day, and third view/tab that displays weather for multiple days. The user **105** may select, for example, the first view/tab, and thus weather data for the present day (which is a portion of the weather data sent by the skill component **190**) may be visible to the user **105**. In such cases, the device **110** is able to provide the presently visible content data, via the widget context data **114**, to the system **120**, which the skill component **190** (and other components of the system **120**) may use for processing.

(54) The skill request router component **188** may send (step **11**) the NLU data corresponding to the user input, the widget context data **116**, the widget identifier, and a command to process with respect to the NLU data to the skill component **190**. The skill component **190** may determine an output responsive to the user input based on the NLU data. In some embodiments, the skill component **190** may determine how the output is to be presented to the user **105**. Such determination may be based on the widget context data **116** and the widget identifier. For example, the skill component **190** may determine that the output is to be presented via the multi-modal widget **112** associated with the widget identifier. As another example, the skill component **190** may determine to launch another multi-modal widget at the device **110** to present the output. As yet another example, the skill component **190** may determine to present the output via a full-screen interface, where the full-screen interface may cause one or more of the multi-modal widgets **112** to appear in the background, or the full-screen interface may be presented on another tab/page of the display screen of the device **110**. As yet another example, the skill component **190** may determine to present the output via an overlay interface that may be presented on-top of the multi-modal widgets **112** and may cover one or more of the multi-modal widgets **112** or cover portions of one or more of the multi-modal widgets **112**.

(55) The skill component **190** may also determine a type of output to be presented to the user **105**. For example, the skill component **190** may determine to present an audio-only output, a visual-only output (including text, image, video, icons, other graphical interface elements, etc.), or an audiovisual output. In some cases, the skill component **190** may determine to output a natural language output. In such cases, the skill component **190** may send text data or SSML tagged data, corresponding to the natural language output, to the orchestrator component **130**. The text data or SSML tagged data may be processed using the TTS component **180** to generate audio data representing synthesized speech, as described in detail below.

(56) The skill component **190** may send (step **12**) output data responsive to the user input to the orchestrator component **130**. The skill component **190** may also send (to the orchestrator component **130**) an indication of how the output data is to be presented to the user **105** and/or the type of output to be presented to the user **105**. The orchestrator component **130** may send (step **13**) the output data to the device **110**, along with a directive/command that causes the device **110** to present the output in the manner specified by the skill component **190**. For example, if the indication from the skill component **190** indicates the output is to be presented via the multi-modal widget **112a**, then, in response to receiving the output data from the orchestrator **130**, the device **110** may update the content of the multi-modal widget **112a**. As another example, if the indication from the skill component **190** indicates the output is to be presented via a full-screen interface, then, in response to receiving the output data from the orchestrator **130**, the device **110** may launch

a full-screen interface to present the output data. The skill component **190**, in some embodiments, may send (at step **12**) an APL document including the output data, and the indication of how the output data is to be presented.

(57) In this manner, the system **120** may determine NLU data and may select a skill component corresponding to a user input using widget context data corresponding to one or more multi-modal widgets being presented at the device **110** that captures the user input.

(58) The widget context data **114** may be provided by the device **110** to the system **120** each time a user input is received and sent to the system **120**. In some cases, the widget context data **114** may remain the same between user inputs, if content at the multi-modal widgets **112** has not updated. In some cases, the widget context data **114** may remain the same after the device **110** is rebooted (restarted) if the content at the multi-modal widgets **112** has not updated.

(59) Each skill component **190** may have multiple multi-modal widgets to provide different or related functionalities. The device **110** may present multiple instances of the same multi-modal widget. In some embodiments, each multi-modal widget **112** may be associated with a widget instance identifier, a widget identifier, and a skill identifier. In some embodiments, the skill selection component **165** may determine the widget instance identifier and the widget identifier corresponding to the user input, and may send the widget instance identifier and the widget identifier to the orchestrator component **130**.

(60) In some embodiments, the widget context data **114**, sent by the device **110** to the system **120**, may include context data corresponding to multi-modal widgets that are not active at the device **110** but are installed or enabled, by the user **105**, at the device **110**.

(61) In some embodiments, a component of the system **120** may request context data to perform its processing. Such context data may include a device identifier for the device **110**, user identifier for the user **105**, profile data for the user **105** from the profile storage **170**, time the user input is received, device location, etc. Such context data may also include widget context data **114** stored at the widget session storage **142**.

(62) In some embodiments, the system **120** may initiate a skill session when a user input from the user **105** is received. The skill session may be used to track data determined during processing of the user input, context data corresponding to the user input, and other data. The skill session may be active while the system **120** is processing with respect to the user input. Once the system **120** presents an output responsive to the user input (e.g., at step **13**) the skill session may be terminated.

(63) In some embodiments, the system **120** may initiate a widget session when the multi-modal widget **112** is enabled at the device **110**. In other embodiments, the system **120** may initiate a widget session when the multi-modal widget **112** is active on the display screen of the device **110**. In some embodiments, the device **110** may send the widget context data **114** when the multi-modal widget **112** becomes enabled or active at the device **110**. The system **120** may store the widget context data **114** in the widget session storage **142**, associated with the device identifier for the device **110** and/or the user identifier for the user **105**. While the multi-modal widget **112** is installed/enabled or is active, the widget session may be active. When the multi-modal widget **112** is disabled or is no longer active at the display screen of the device **110**, the widget session may be terminated. The widget session may be used to track widget context data corresponding to the multi-modal widget **112**. The widget context data may be updated/refreshed when the multi-modal widget **112** is updated at the device **110** (e.g., with respect to content, display position, receipt of user input, etc.). Other components of the system **120** may request the widget context data to perform processing. In some cases, a lifecycle of a widget session may be longer than a skill session, as the skill session may terminate when an interaction ends (e.g., the system presents an output responsive to a user input), while the widget session may terminate when the multi-modal widget is disabled or no longer active at the device **110**.

(64) In some embodiments, the data stored at the widget session storage **142** (and corresponding to historic interactions between the user **105** and the multi-modal widgets **112**) may be used by the

NLU component **160** and the skill selection component **165** to determine how to respond to a user input corresponding to a multi-modal widget. For example, the user **105** may say “save that recipe” intending to save a recipe being displayed via the multi-modal widget **112a** at the device **110**. Based on processing the user input as described herein, the system **120** performs the action of saving the recipe displayed at the multi-modal widget **112a**. Subsequently (after some time after saying “save that recipe”) the user **105** may say “unsave that recipe” or “delete that recipe, I changed my mind” or other similar inputs. In this case, the NLU component **160** may determine a “reverse action” or “undo action” intent corresponding to the subsequent user input. Using the data stored at the widget session storage **142**, the NLU component **160** may determine that the intent corresponds to the multi-modal widget **112a** based on the historic interaction data indicating that a recipe was previously saved with respect to the multi-modal widget **112a**. The skill selection component **165** may select the skill component **190**, corresponding to the multi-modal widget **112a**, to respond to the subsequent user input (e.g., perform the action of deleting the recipe).

(65) The system **120** may use other components illustrated in FIG. **1**. The various components shown in FIG. **1** may be located on a same or different physical devices. Communication between various components may occur directly or across the network(s) **199**.

(66) A microphone or array of microphones (of or otherwise associated with the device **110**) may capture audio. The device **110** processes audio data, representing the audio, to determine whether speech is detected. The device **110** may use various techniques to determine whether audio data includes speech. In some examples, the device **110** may apply voice activity detection (VAD) techniques. Such techniques may determine whether speech is present in audio data based on various quantitative aspects of the audio data, such as the spectral slope between one or more frames of the audio data, the energy levels of the audio data in one or more spectral bands, the signal-to-noise ratios of the audio data in one or more spectral bands, or other quantitative aspects. In other examples, the device **110** may implement a classifier configured to distinguish speech from background noise. The classifier may be implemented by techniques such as linear classifiers, support vector machines, and decision trees. In still other examples, the device **110** may apply Hidden Markov Model (HMM) or Gaussian Mixture Model (GMM) techniques to compare the audio data to one or more acoustic models in storage, which acoustic models may include models corresponding to speech, noise (e.g., environmental noise or background noise), or silence. Still other techniques may be used to determine whether speech is present in audio data.

(67) Once speech is detected in audio data, the device **110** may determine if the speech is directed at the device **110**/system **120**. In at least some embodiments, such determination may be made using a wakeword detection component **320** (shown in FIG. **3**). The wakeword detection component **320** may be configured to detect various wakewords. In at least some examples, each wakeword may correspond to a name of a different digital assistant. An example wakeword/digital assistant name is “Alexa.”

(68) Wakeword detection is typically performed without performing linguistic analysis, textual analysis, or semantic analysis. Instead, the audio data is analyzed to determine if specific characteristics of the audio data match preconfigured acoustic waveforms, audio signatures, or other data corresponding to a wakeword.

(69) Thus, the wakeword detection component **320** may compare audio data to stored data to detect a wakeword. One approach for wakeword detection applies general large vocabulary continuous speech recognition (LVCSR) systems to decode audio signals, with wakeword searching being conducted in the resulting lattices or confusion networks. Another approach for wakeword detection builds HMMs for each wakeword and non-wakeword speech signals, respectively. The non-wakeword speech includes other spoken words, background noise, etc. There can be one or more HMMs built to model the non-wakeword speech characteristics, which are named filler models. Viterbi decoding is used to search the best path in the decoding graph, and the decoding output is further processed to make the decision on wakeword presence. This approach can be

extended to include discriminative information by incorporating a hybrid deep neural network (DNN)-HMM decoding framework. In another example, the wakeword detection component **320** may be built on DNN/recursive neural network (RNN) structures directly, without HMM being involved. Such an architecture may estimate the posteriors of wakewords with context data, either by stacking frames within a context window for DNN, or using RNN. Follow-on posterior threshold tuning or smoothing is applied for decision making. Other techniques for wakeword detection, such as those known in the art, may also be used.

(70) Once the wakeword detection component **320** detects a wakeword, the device **110** may “wake” and begin transmitting audio data **311**, representing the audio, to the system **120**. The audio data **311** may include the detected wakeword, or the device **110** may remove the portion of the audio data, corresponding to the detected wakeword, prior to sending the audio data **311** to the system **120**.

(71) Referring to FIG. **1**, the orchestrator component **130** may be configured to, among other things, coordinate data transmissions between components of the system **120**. The orchestrator component **130** may receive audio data from the device **110**, and send the audio data to the ASR component **150**.

(72) The ASR component **150** transcribes the audio data into ASR output data including one or more ASR hypotheses. An ASR hypothesis may be configured as a textual interpretation of the speech in the audio data, or may be configured in another manner, such as one or more tokens. Each ASR hypothesis may represent a different likely interpretation of the speech in the audio data. Each ASR hypothesis may be associated with a score (e.g., confidence score, probability score, or the like) representing the associated ASR hypothesis correctly represents the speech in the audio data.

(73) The ASR component **150** interprets the speech in the audio data based on a similarity between the audio data and pre-established language models. For example, the ASR component **150** may compare the audio data **311** with models for sounds (e.g., subword units, such as phonemes, etc.) and sequences of sounds to identify words that match the sequence of sounds of the speech represented in the audio data.

(74) In at least some instances, instead of the device **110** receiving a spoken natural language input, the device **110** may receive a textual (e.g., types) natural language input. The device **110** may determine text data representing the textual natural language input, and may send the text data to the system **120**, wherein the text data is received by the orchestrator component **130**. The orchestrator component **130** may send the text data or ASR output data, depending on the type of natural language input received, to the NLU component **160**.

(75) The NLU component **160** processes the ASR output data or text data to determine one or more NLU hypotheses embodied in NLU output data. The NLU component **160** may perform intent classification (IC) processing on the ASR output data or text data to determine an intent of the natural language input. An intent corresponds to an action to be performed that is responsive to the natural language input. To perform IC processing, the NLU component **160** may communicate with a database of words linked to intents. For example, a music intent database may link words and phrases such as “quiet,” “volume off,” and “mute” to a <Mute> intent. The NLU component **160** identifies intents by comparing words and phrases in ASR output data or text data to the words and phrases in an intents database. In some embodiments, the NLU component **160** may communicate with multiple intents databases, with each intents database corresponding to one or more intents associated with a particular skill.

(76) For example, IC processing of the natural language input “play my workout playlist” may determine an intent of <PlayMusic>. For further example, IC processing of the natural language input “call mom” may determine an intent of <Call>. In another example, IC processing of the natural language input “call mom using video” may determine an intent of <VideoCall>. In yet another example, IC processing of the natural language input “what is today's weather” may

determine an intent of <OutputWeather>.

(77) The NLU component **160** may also perform named entity recognition (NER) processing on the ASR output data or text data to determine one or more portions, sometimes referred to as slots, of the natural language input that may be needed for post-NLU processing (e.g., processing performed by a skill). For example, NER processing of the natural language input “play [song name]” may determine an entity type of “SongName” and an entity value corresponding to the indicated song name. For further example, NER processing of the natural language input “call mom” may determine an entity type of “Recipient” and an entity value corresponding to “mom.” In another example, NER processing of the natural language input “what is today's weather” may determine an entity type of “Date” and an entity value of “today.”

(78) In at least some embodiments, the intents identifiable by the NLU component **160** may be linked to one or more grammar frameworks with entity types to be populated with entity values. Each entity type of a grammar framework corresponds to a portion of ASR output data or text data that the NLU component **160** believes corresponds to an entity value. For example, a grammar framework corresponding to a <PlayMusic> intent may correspond to sentence structures such as “Play {Artist Name},” “Play {Album Name},” “Play {Song name},” “Play {Song name} by {Artist Name},” etc.

(79) For example, the NLU component **160** may perform NER processing to identify words in ASR output data or text data as subject, object, verb, preposition, etc. based on grammar rules and/or models. Then, the NLU component **160** may perform IC processing using the identified verb to identify an intent. Thereafter, the NLU component **160** may again perform NER processing to determine a grammar model associated with the identified intent. For example, a grammar model for a <PlayMusic> intent may specify a list of entity types applicable to play the identified “object” and any object modifier (e.g., a prepositional phrase), such as {Artist Name}, {Album Name}, {Song name}, etc. The NER processing may then involve searching corresponding fields in a lexicon, attempting to match words and phrases in the ASR output data that NER processing previously tagged as a grammatical object or object modifier with those identified in the lexicon.

(80) NER processing may include semantic tagging, which is the labeling of a word or combination of words according to their type/semantic meaning. NER processing may include parsing ASR output data or text data using heuristic grammar rules, or a model may be constructed using techniques such as hidden Markov models, maximum entropy models, log linear models, conditional random fields (CRFs), and the like. For example, NER processing with respect to a music skill may include parsing and tagging ASR output data or text data corresponding to “play mother's little helper by the rolling stones” as {Verb}: “Play,” {Object}: “mother's little helper,” {Object Preposition}: “by,” and {Object Modifier}: “the rolling stones.” The NER processing may identify “Play” as a verb based on a word database associated with the music skill, which IC processing determines corresponds to a <PlayMusic> intent.

(81) The NLU component **160** may generate NLU output data including one or more NLU hypotheses, with each NLU hypothesis including an intent and optionally one or more entity types and corresponding entity values. In some embodiments, the NLU component **160** may perform IC processing and NER processing with respect to different skills. One skill may support the same or different intents than another skill. Thus, the NLU output data may include multiple NLU hypotheses, with each NLU hypothesis corresponding to IC processing and NER processing performed on the ASR output or text data with respect to a different skill.

(82) As described above, the system **120** may perform speech processing using two different components (e.g., the ASR component **150** and the NLU component **160**). In at least some embodiments, the system **120** may implement a spoken language understanding (SLU) component configured to process audio data **311** to determine NLU output data.

(83) The SLU component may be equivalent to a combination of the ASR component **150** and the NLU component **160**. Yet, the SLU component may process audio data **311** and directly determine

the NLU output data, without an intermediate step of generating ASR output data. As such, the SLU component may take audio data **311** representing a spoken natural language input and attempt to make a semantic interpretation of the spoken natural language input. That is, the SLU component may determine a meaning associated with the spoken natural language input and then implement that meaning. For example, the SLU component may interpret audio data **311** representing a spoken natural language input in order to derive a desired action. The SLU component may output a most likely NLU hypothesis, or multiple NLU hypotheses associated with respective confidence or other scores (such as probability scores, etc.).

(84) The system **120** may include one or more skill components **190** and/or may communicate with one or more skills **125**. A “skill” may refer to software, that may be placed on a machine or a virtual machine (e.g., software that may be launched in a virtual instance when called), configured to process NLU output data and perform one or more actions in response thereto. For example, for NLU output data including a <PlayMusic> intent, an “artist” entity type, and an artist name as an entity value, a music skill may be called to output music sung by the indicated artist. For further example, for NLU output data including a <TurnOn> intent, a “device” entity type, and an entity value of “lights,” a smart home skill may be called to cause one or more “smart” lights to operate in an “on” state. In another example, for NLU output data including an <OutputWeather> intent, a “location” entity type, and an entity value corresponding to a geographic location of the device **110**, a weather skill may be called to output weather information for the geographic location. For further example, for NLU output data including a <BookRide> intent, a taxi skill may be called to book a requested ride. In another example, for NLU output data including a <BuyPizza> intent, a restaurant skill may be called to place an order for a pizza.

(85) A skill component **190** may operate in conjunction between the system **120** and other devices, such as the device **110**, a restaurant electronic ordering system, a taxi electronic booking system, etc. in order to complete certain functions. Inputs to a skill component **190** may come from speech processing interactions or through other interactions or input sources.

(86) A skill component **190** may be associated with a domain, a non-limiting list of which includes a smart home domain, a music domain, a video domain, a weather domain, a communications domain, a flash briefing domain, a shopping domain, and a custom domain.

(87) The system **120** may include a TTS component **180** that generates audio data including synthesized speech. The data input to the TTS component **180** may come from a skill component **190**, the orchestrator component **130**, or another component of the system **120**. The data input to the TTS component **180** may be text data or synthesized speech markup language (SSML) tagged data.

(88) In one method of synthesis called unit selection, the TTS component **180** matches input data against a database of recorded speech. The TTS component **180** selects matching units of recorded speech and concatenates the units together to form audio data. In another method of synthesis called parametric synthesis, the TTS component **180** varies parameters such as frequency, volume, and noise to determine audio data including an artificial speech waveform. Parametric synthesis uses a computerized voice generator, sometimes called a vocoder.

(89) The system **120** may include a user recognition component **195**. The user recognition component **195** may recognize one or more users using various data. The user recognition component **195** may take as input the audio data **311**. The user recognition component **195** may perform user recognition by comparing speech characteristics, in the audio data **311**, to stored speech characteristics of users. The user recognition component **195** may additionally or alternatively perform user recognition by comparing biometric data (e.g., fingerprint data, iris data, retina data, etc.), received by the system **120** in correlation with a natural language input, to stored biometric data of users. The user recognition component **195** may additionally or alternatively perform user recognition by comparing image data (e.g., including a representation of at least a feature of a user), received by the system **120** in correlation with a natural language input, with

stored image data including representations of features of different users. The user recognition component **195** may perform other or additional user recognition processes, including those known in the art. For a particular natural language input, the user recognition component **195** may perform processing with respect to stored data of users associated with the device **110** that received the natural language input.

(90) The user recognition component **195** determines whether a natural language input originated from a particular user. For example, the user recognition component **195** may determine a first value representing a likelihood that a natural language input originated from a first user, a second value representing a likelihood that the natural language input originated from a second user, etc. The user recognition component **195** may also determine an overall confidence regarding the accuracy of user recognition processing.

(91) The user recognition component **195** may output a single user identifier corresponding to the most likely user that originated the natural language input. Alternatively, the user recognition component **195** may output multiple user identifiers (e.g., in the form of an N-best list) with respective values representing likelihoods of respective users originating the natural language input. The output of the user recognition component **195** may be used to inform NLU processing, processing performed by a skill **125**, as well as processing performed by other components of the system **120** and/or other systems.

(92) The system **120** may include profile storage **170**. The profile storage **170** may include a variety of data related to individual users, groups of users, devices, etc. that interact with the system **120**. As used herein, a “profile” refers to a set of data associated with a user, group of users, device, etc. The data of a profile may include preferences specific to the user, group of users, device, etc.; input and output capabilities of one or more devices; internet connectivity data; user bibliographic data; subscription data; skill enablement data; and/or other data.

(93) The profile storage **170** may include one or more user profiles. Each user profile may be associated with a different user identifier. Each user profile may include various user identifying data (e.g., name, gender, address, language(s), etc.). Each user profile may also include preferences of the user. Each user profile may include one or more device identifiers, each representing a respective device registered to the user. Each user profile may include skill identifiers of skills **125** that the user has enabled. When a user enables a skill **125**, the user is providing the system **120** with permission to allow the skill **125** to execute with respect to the user's natural language inputs. If a user does not enable a skill **125**, the system **120** may not execute the skill **125** with respect to the user's natural language inputs.

(94) The profile storage **170** may include one or more group profiles. Each group profile may be associated with a different group identifier. A group profile may be specific to a group of users. That is, a group profile may be associated with two or more individual user profiles. For example, a group profile may be a household profile that is associated with user profiles associated with multiple users of a single household. A group profile may include preferences shared by all the user profiles associated therewith. Each user profile associated with a group profile may additionally include preferences specific to the user associated therewith. That is, a user profile may include preferences unique from one or more other user profiles associated with the same group profile. A user profile may be a stand-alone profile or may be associated with a group profile. A group profile may be associated with (or include) one or more device profiles corresponding to one or more devices associated with the group profile.

(95) The profile storage **170** may include one or more device profiles. Each device profile may be associated with a different device identifier. A device profile may include various device identifying data, input/output characteristics, networking characteristics, etc. A device profile may also include one or more user identifiers, corresponding to one or more user profiles associated with the device profile. For example, a household device's profile may include the user identifiers of users of the household.

(96) The foregoing describes illustrative components and processing of the system **120**. The following describes illustrative components and processing of the device **110**. As illustrated in FIG. **3**, in at least some embodiments the system **120** may receive audio data **311** from the device **110**, to recognize speech corresponding to a spoken natural language in the received audio data **311**, and to perform functions in response to the recognized speech. In at least some embodiments, these functions involve sending directives (e.g., commands), from the system **120** to the device **110** to cause the device **110** to perform an action, such as output synthesized speech (responsive to the spoken natural language input) via a loudspeaker(s), and/or control one or more secondary devices by sending control commands to the one or more secondary devices.

(97) Thus, when the device **110** is able to communicate with the system **120** over the network(s) **199**, some or all of the functions capable of being performed by the system **120** may be performed by sending one or more directives over the network(s) **199** to the device **110**, which, in turn, may process the directive(s) and perform one or more corresponding actions. For example, the system **120**, using a remote directive that is included in response data (e.g., a remote response), may instruct the device **110** to output synthesized speech via a loudspeaker(s) of (or otherwise associated with) the device **110**, to output content (e.g., music) via the loudspeaker(s) of (or otherwise associated with) the device **110**, to display content on a display of (or otherwise associated with) the device **110**, and/or to send a directive to a secondary device (e.g., a directive to turn on a smart light). It will be appreciated that the system **120** may be configured to provide other functions in addition to those discussed herein, such as, without limitation, providing step-by-step directions for navigating from an origin location to a destination location, conducting an electronic commerce transaction on behalf of the user **105** as part of a shopping function, establishing a communication session (e.g., an audio or video call) between the user **105** and another user, and so on.

(98) The device **110** may include a wakeword detection component **320** configured to detect a wakeword (e.g., “Alexa”) that indicates to the device **110** that the audio data **311** is to be processed for determining NLU output data. In at least some embodiments, a hybrid selector **324**, of the device **110**, may send the audio data **311** to the wakeword detection component **320**. If the wakeword detection component **320** detects a wakeword in the audio data **311**, the wakeword detection component **320** may send an indication of such detection to the hybrid selector **324**. In response to receiving the indication, the hybrid selector **324** may send the audio data **311** to the system **120** and/or an on-device ASR component **150**. The wakeword detection component **320** may also send an indication, to the hybrid selector **324**, representing a wakeword was not detected. In response to receiving such an indication, the hybrid selector **324** may refrain from sending the audio data **311** to the system **120**, and may prevent the on-device ASR component **150** from processing the audio data **311**. In this situation, the audio data **311** can be discarded.

(99) The device **110** may conduct its own speech processing using on-device language processing components (such as an on-device SLU component, an on-device ASR component **350**, and/or an on-device NLU component **360**) similar to the manner discussed above with respect to the speech processing system-implemented ASR component **150**, and NLU component **160**. The device **110** may also internally include, or otherwise have access to, other components such as one or more skills **190**, a user recognition component **395** (configured to process in a similar manner to the user recognition component **195**), profile storage **370** (configured to store similar profile data to the profile storage **170**), a TTS component **380** (configured to process in a similar manner to the TTS component **180**), one or more skill component(s) **390** (configured to process in a similar manner to the skill component(s) **190**), and other components. In at least some embodiments, the on-device profile storage **370** may only store profile data for a user or group of users specifically associated with the device **110**.

(100) In at least some embodiments, the on-device language processing components may not have the same capabilities as the language processing components implemented by the system **120**. For

example, the on-device language processing components may be configured to handle only a subset of the natural language inputs that may be handled by the speech processing system-implemented language processing components. For example, such subset of natural language inputs may correspond to local-type natural language inputs, such as those controlling devices or components associated with a user's home. In such circumstances the on-device language processing components may be able to more quickly interpret and respond to a local-type natural language input, for example, than processing that involves the system **120**. If the device **110** attempts to process a natural language input for which the on-device language processing components are not necessarily best suited, the NLU output data, determined by the on-device components, may have a low confidence or other metric indicating that the processing by the on-device language processing components may not be as accurate as the processing done by the system **120**.

(101) The hybrid selector **324**, of the device **110**, may include a hybrid proxy (HP) **326** configured to proxy traffic to/from the system **120**. For example, the HP **326** may be configured to send messages to/from a hybrid execution controller (HEC) **327** of the hybrid selector **324**. For example, command/directive data received from the system **120** can be sent to the HEC **327** using the HP **326**. The HP **326** may also be configured to allow the audio data **311** to pass to the system **120** while also receiving (e.g., intercepting) this audio data **311** and sending the audio data **311** to the HEC **327**.

(102) In at least some embodiments, the hybrid selector **324** may further include a local request orchestrator (LRO) **328** configured to notify the on-device ASR component **150** about the availability of the audio data **311**, and to otherwise initiate the operations of on-device language processing when the audio data **311** becomes available. In general, the hybrid selector **324** may control execution of on-device language processing, such as by sending “execute” and “terminate” events/instructions. An “execute” event may instruct a component to continue any suspended execution (e.g., by instructing the component to execute on a previously-determined intent in order to determine a directive). Meanwhile, a “terminate” event may instruct a component to terminate further execution, such as when the device **110** receives directive data from the system **120** and chooses to use that remotely-determined directive data.

(103) Thus, when the audio data **311** is received, the HP **326** may allow the audio data **311** to pass through to the system **120** and the HP **326** may also input the audio data **311** to the on-device ASR component **150** by routing the audio data **311** through the HEC **327** of the hybrid selector **324**, whereby the LRO **328** notifies the on-device ASR component **150** of the audio data **311**. At this point, the hybrid selector **324** may wait for response data from either or both the system **120** and/or the on-device language processing components. However, the disclosure is not limited thereto, and in some examples the hybrid selector **324** may send the audio data **311** only to the on-device ASR component **150** without departing from the disclosure. For example, the device **110** may process the audio data **311** on-device without sending the audio data **311** to the system **120**.

(104) The on-device ASR component **350** is configured to receive the audio data **311** from the hybrid selector **324**, and to recognize speech in the audio data **311**, and the on-device NLU component **360** is configured to determine an intent from the recognized speech (an optionally one or more named entities), and to determine how to act on the intent by generating NLU output data that may include directive data (e.g., instructing a component to perform an action). In some cases, a directive may include a description of the intent (e.g., an intent to turn off {device A}). In some cases, a directive may include (e.g., encode) an identifier of a second device(s), such as kitchen lights, and an operation to be performed at the second device(s). Directive data may be formatted using Java, such as JavaScript syntax, or JavaScript-based syntax. This may include formatting the directive using JSON. In at least some embodiments, a device-determined directive may be serialized, much like how remotely-determined directives may be serialized for transmission in data packets over the network(s) **199**. In at least some embodiments, a device-determined directive may be formatted as a programmatic application programming interface (API) call with a same

logical operation as a remotely-determined directive. In other words, a device-determined directive may mimic a remotely-determined directive by using a same, or a similar, format as the remotely-determined directive.

(105) A NLU hypothesis (output by the on-device NLU component **360**) may be selected as usable to respond to a natural language input, and local response data may be sent (e.g., local NLU output data, local knowledge base information, internet search results, and/or local directive data) to the hybrid selector **324**, such as a “ReadyToExecute” response. The hybrid selector **324** may then determine whether to use directive data from the on-device components to respond to the natural language input, to use directive data received from the system **120**, assuming a remote response is even received (e.g., when the device **110** is able to access the system **120** over the network(s) **199**), or to determine output data requesting additional information from the user **105**.

(106) The device **110** and/or the system **120** may associate a unique identifier with each natural language input. The device **110** may include the unique identifier when sending the audio data **311** to the system **120**, and the response data from the system **120** may include the unique identifier to identify to which natural language input the response data corresponds.

(107) In at least some embodiments, the device **110** may include one or more skill components **190**. The skill component(s) **190** installed on (or in communication with) the device **110** may include, without limitation, a smart home skill and/or a device control skill configured to control a second device(s), a music skill configured to output music, a navigation skill configured to output directions, a shopping skill configured to conduct an electronic purchase, and/or the like.

(108) One or more of the components described herein may employ a machine learning (ML) model(s). Generally, ML models may be trained and operated according to various machine learning techniques. Such techniques may include, for example, neural networks (such as deep neural networks and/or recurrent neural networks), inference engines, trained classifiers, etc. Examples of trained classifiers include Support Vector Machines (SVMs), neural networks, decision trees, AdaBoost (short for “Adaptive Boosting”) combined with decision trees, and random forests. Focusing on SVM as an example, SVM is a supervised learning model with associated learning algorithms that analyze data and recognition patterns in the data, and which are commonly used for classification and regression analysis. Given a set of training examples, each marked as belonging to one of two categories, an SVM training algorithm builds a model that assigns new examples into one category or the other, making it a non-probabilistic binary linear classifier. More complex SVM models may be built with the training set identifying more than two categories, with the SVM determining which category is most similar to input data. An SVM model may be mapped so that the examples of the separate categories are divided by clear gaps. New examples are then mapped into that same space and predicted to belong to a category based on which side of the gaps they fall on. Classifiers may issue a “score” indicating which category the data most closely matches. The score may provide an indication of how closely the data matches the category.

(109) In order to apply machine learning techniques, machine learning processes themselves need to be trained. Training a machine learning model requires establishing a “ground truth” for the training examples. In machine learning, the term “ground truth” refers to the accuracy of a training set's classification for supervised learning techniques. Various techniques may be used to train the models including backpropagation, statistical learning, supervised learning, semi-supervised learning, stochastic learning, or other known techniques.

(110) As used herein, a “domain” may refer to a collection of related functionality. A domain may be associated with one or more skills performing related functionality may be part of a domain. A non-limiting list of domains includes a smart home domain (corresponding to smart home functionality), a music domain (corresponding to music functionality), a video domain (corresponding to video functionality), a weather domain (corresponding to weather functionality), a communications domain (corresponding to one- or two-way communications functionality), and a

shopping domain (corresponding to shopping functionality).

(111) Referring now to FIG. 4, it is described how NLU processing may be performed. The NLU component **160/360** may include one or more recognizers **463**. In at least some embodiments, a recognizer **463** may be associated with a skill component **190/390** (e.g., the recognizer may be configured to interpret a natural language input to correspond to the skill component **190/390**). In at least some other examples, a recognizer **463** may be associated with a domain (e.g., the recognizer may be configured to interpret a natural language input to correspond to the domain). A non-limiting list of domains includes a smart home domain, a music domain, a video domain, a flash briefing domain, a shopping domain, a communications domain, and/or a custom domain.

(112) Recognizers **463** may process in parallel, in series, partially in parallel, etc. For example, a recognizer corresponding to a first domain may process at least partially in parallel to a recognizer corresponding to a second domain. For further example, a recognizer corresponding to a first skill component may process at least partially in parallel to a recognizer corresponding to a second skill component. In another example, a recognizer corresponding to a domain may process at least partially in parallel to a recognizer corresponding to a skill component.

(113) The NLU component **160/360** may communicate with various storages. The NLU component **160/360** may communicate with an NLU storage, which includes skill component grammars, representing how natural language inputs may be formulated to invoke skill components **190/390**, and skill component intents representing intents supported by respective skill components **190/390**.

(114) Each recognizer **463** may be associated with a particular grammar, one or more particular intents, and a particular personalized lexicon (stored in an entity library). A gazetteer may include skill component-indexed lexical information associated with a particular user. For example, Gazetteer A may include skill component-indexed lexical information. A user's music skill component lexical information might include album titles, artist names, and song names, for example, whereas a user's contact list skill component lexical information might include the names of contacts. Since every user's music collection and contact list is presumably different, this personalized information may improve entity resolution.

(115) Each recognizer **463** may include a named entity recognition (NER) component **462** that attempts to identify grammars and lexical information that may be used to construe meaning with respect to text data input therein. A NER component **462** identifies portions of text data that correspond to a named entity that may be recognizable by the system **120/device 110**. A NER component **462** may also determine whether a word refers to an entity that is not explicitly mentioned in the text, for example “him,” “her,” “it” or other anaphora, exophora or the like.

(116) A NER component **462** applies grammar models and lexical information associated with one or more skill components **190/390** to determine a mention of one or more entities in text data input therein. In this manner, a NER component **462** identifies “slots” (i.e., particular words in text data) that may be needed for later processing. A NER component **462** may also label each slot with a type (e.g., noun, place, city, artist name, song name, etc.).

(117) Each grammar model may include the names of entities (i.e., nouns) commonly found in natural language about a particular skill component **190/390** to which the grammar model relates, whereas lexical information may be personalized to the user identifier output by a user recognition component **190/390** for the natural language input. For example, a grammar model associated with a shopping skill component may include a database of words commonly used when people discuss shopping.

(118) A downstream process called named entity resolution links a portion of input data (identified by a NER component **462**) to a specific entity known to the system **120/device 110**. To perform named entity resolution, the NLU component **160/360** may use gazetteer information stored in the entity library storage. The gazetteer information may be used to match text data (identified by a NER component **462**) with different entities, such as song titles, contact names, etc. Gazetteers may be linked to users (e.g., a particular gazetteer may be associated with a specific user's music

collection), may be linked to certain skill components **190/390** (e.g., a shopping skill component, a music skill component, a video skill component, a communications skill component, etc.), or may be organized in another manner.

(119) Each recognizer **463** may also include an intent classification (IC) component **464** that processes text data input thereto to determine an intent(s) of a skill component(s) **190/390** that potentially corresponds to the natural language input represented in the text data. An intent corresponds to an action to be performed that is responsive to the natural language input represented by the text data. An IC component **464** may communicate with a database of words linked to intents. For example, a music intent database may link words and phrases such as “quiet,” “volume off,” and “mute” to a <Mute> intent. An IC component **464** identifies potential intents by comparing words and phrases in text data to the words and phrases in an intents database associated with the skill component(s) **190/390** that is associated with the recognizer **463** implementing the IC component **464**.

(120) The intents identifiable by a specific IC component **464** may be linked to one or more skill component-specific grammar frameworks with “slots” to be filled. Each slot of a grammar framework corresponds to a portion of text data that a NER component **462** believes corresponds to an entity. For example, a grammar framework corresponding to a <PlayMusic> intent may correspond to text data sentence structures such as “Play {Artist Name},” “Play {Album Name},” “Play {Song name},” “Play {Song name} by {Artist Name},” etc. However, to make resolution more flexible, grammar frameworks may not be structured as sentences, but rather based on associating slots with grammatical tags.

(121) For example, a NER component **462** may identify words in text data as subject, object, verb, preposition, etc. based on grammar rules and/or models prior to recognizing named entities in the text data. An IC component **464** (implemented by the same recognizer **463**) may use the identified verb to identify an intent. The NER component **462** may then determine a grammar model associated with the identified intent. For example, a grammar model for an intent corresponding to <PlayMusic> may specify a list of slots applicable to play the identified “object” and any object modifier (e.g., a prepositional phrase), such as {Artist Name}, {Album Name}, {Song name}, etc. The NER component **462** may then search corresponding fields in a lexicon, attempting to match words and phrases in the text data the NER component **462** previously tagged as a grammatical object or object modifier with those identified in the lexicon.

(122) A NER component **462** may perform semantic tagging, which is the labeling of a word or combination of words according to their type/semantic meaning. A NER component **462** may parse text data using heuristic grammar rules, or a model may be constructed using techniques such as hidden Markov models, maximum entropy models, log linear models, conditional random fields (CRF), and the like. For example, a NER component **462**, implemented by a music skill component or music domain recognizer **463**, may parse and tag text data corresponding to “play mother's little helper by the rolling stones” as {Verb}: “Play,” {Object}: “mother's little helper,” {Object Preposition}: “by,” and {Object Modifier}: “the rolling stones.” The NER component **462** may identify “Play” as a verb based on a word database associated with the music skill component or music domain, which an IC component **464** may determine corresponds to a <PlayMusic> intent. At this stage, no determination has been made as to the meaning of “mother's little helper” and “the rolling stones,” but based on grammar rules and models, the NER component **462** has determined that the text of these phrases relates to the grammatical object (i.e., entity).

(123) The frameworks linked to the intent are then used to determine what database fields should be searched to determine the meaning of these phrases, such as searching a user's gazetteer for similarity with the framework slots. For example, a framework for a <PlayMusic> intent might indicate to attempt to resolve the identified object based on {Artist Name}, {Album Name}, and {Song name}, and another framework for the same intent might indicate to attempt to resolve the object modifier based on {Artist Name}, and resolve the object based on {Album Name} and

{Song Name} linked to the identified {Artist Name}. If the search of the gazetteer does not resolve a slot/field using gazetteer information, the NER component **462** may search a database of generic words (in a knowledge base). For example, if the text data corresponds to “play songs by the rolling stones,” after failing to determine an album name or song name called “songs” by “the rolling stones,” the NER component **462** may search a music skill component vocabulary for the word “songs.” In the alternative, generic words may be checked before the gazetteer information, or both may be tried, potentially producing two different results.

(124) A recognizer **463** may tag text data to attribute meaning thereto. For example, a recognizer **463** may tag “play mother's little helper by the rolling stones” as: {skill component} Music, {intent} Play Music, {artist name} rolling stones, {media type} SONG, and {song title} mother's little helper. For further example, a recognizer **463** may tag “play songs by the rolling stones” as: {skill component} Music, {intent} Play Music, {artist name} rolling stones, and {media type} SONG.

(125) As described above, more than one recognizer **463** may process with respect to ASR data **402** representing a single natural language input. In such instances, each recognizer **463** may output at least one NLU hypothesis including an intent indicator (determined by an IC component **464** of the recognizer **463**) and at least one tagged named entity (determined by a NER component **462** of the recognizer **463**).

(126) In some embodiments, the recognizer **463**, including the NER component **462** and the IC component **464**, may process with respect to the widget context data **114**. The IC component **464** may use the information included in the widget context data **114** to determine intent data corresponding to a user input. The NER component **462** may use the information included in the widget context data **114** to determine entity data.

(127) The NLU component **160/360** may include a shortlister component **410**. The shortlister component **410** selects skill components **190/390** that may execute in response to the natural language input. The shortlister component **410** thus limits downstream, more resource intensive NLU processes to being performed with respect to skill components **190/390** that are likely to execute in response to the natural language input.

(128) Without a shortlister component **410**, the NLU component **160/360** may process a given ASR hypothesis, included in the ASR data **402**, (or the text data depending on the type of natural language input being processed) with respect to every skill component **190/390** of (or in communication with) the system **120**, either in parallel, in series, or using some combination thereof. By implementing a shortlister component **410**, the NLU component **160/360** may process a given ASR hypothesis included in the ASR data **402** with respect to only the skill components **190/390** that are likely to execute in response to the natural language input. This reduces total compute power and latency attributed to NLU processing.

(129) The shortlister component **410** may include one or more machine learning (ML) models. The ML model(s) may be trained to recognize various forms of natural language inputs that may be received by the system **120/device 110**. For example, during a training period a skill component **190/390** may provide the system **120/device 110** with training data representing sample natural language inputs that may be provided by a user to invoke the skill component **190/390**. For example, a ride sharing skill component may provide the system **120/device 110** with training data including text corresponding to “get me a cab to [location],” “get me a ride to [location],” “book me a cab to [location],” “book me a ride to [location],” etc. The one or more ML models may be trained using the training data to determine other potentially related natural language input structures that a user may try to use to invoke the particular skill component **190/390**. During training, the system **120/device 110** may query the skill component **190/390** regarding whether the determined other natural language input structures are permissible, from the perspective of the skill component **190/390**, to be used to invoke the skill component **190/390** at runtime. The alternate natural language input structures may be derived by one or more ML models during model training

and/or may be based on natural language input structures provided by different skill components **190/390**. The skill component **190/390** may also provide the system **120/device 110** with training data indicating grammar and annotations. The system **120/device 110** may use the training data representing the sample natural language inputs, the determined related natural language input(s), the grammar, and the annotations to train a ML model that indicates when a runtime natural language input is likely to be directed to/handled by a particular skill component **190/390**. Each ML model of the shortlister component **410** may be trained with respect to a different skill component **190/390**. Alternatively, the shortlister component **410** may implement one ML model per skill component type, such as one ML model for weather skill components, one ML model for ride sharing skill components, etc.

(130) The system **120/device 110** may use the sample natural language inputs provided by a skill component **190/390**, and related sample natural language inputs determined during training, as binary examples to train a ML model associated with the skill component **190/390**. The ML model associated with the particular skill component **190/390** may then be operated at runtime by the shortlister component **410**. Some sample natural language inputs may be positive training examples (e.g., natural language inputs that may be used to invoke the skill component **190/390**), whereas other sample natural language inputs may be negative training examples (e.g., natural language inputs that may not be used to invoke the skill component **190/390**).

(131) As described above, the shortlister component **410** may include a different ML model for each skill component **190/390**, a different ML model for each skill component type, or some other combination of ML models. For example, the shortlister component **410** may alternatively include a single ML model. The single ML model may include a portion trained with respect to characteristics (e.g., semantic characteristics) shared by all skill components **190/390**. The single ML model may also include skill component-specific portions, with each skill component-specific portion being trained with respect to a different skill component **190/390**. Implementing a single ML model with skill component-specific portions may result in less latency than implementing a different ML model for each skill component **190/390** because the single ML model with skill component-specific portions limits the number of characteristics processed on a per skill component level.

(132) The portion of the ML model, trained with respect to characteristics shared by more than one skill component **190/390**, may be clustered based on skill component type. For example, a first portion, of the portion trained with respect to multiple skill components **190/390**, may be trained with respect to weather skill components; a second portion, of the portion trained with respect to multiple skill components **190/390**, may be trained with respect to music skill components; a third portion, of the portion trained with respect to multiple skill components **190/390**, may be trained with respect to travel skill components; etc.

(133) Clustering may not be beneficial in every instance because clustering may cause the shortlister component **410** to output indications of only a portion of the skill components **190/390** that the natural language input may relate to. For example, a natural language input may correspond to “tell me about Tom Collins.” If the ML model is clustered based on skill component type, the shortlister component **410** may determine the natural language input corresponds to a recipe skill component (e.g., storing or otherwise having access to a drink recipe) even though the natural language input may also correspond to an information skill component (e.g., storing or otherwise having access to information about a person named Tom Collins).

(134) Training the ML model(s) of the shortlister component **410** may require establishing a “ground truth” for the training examples input therein. In machine learning, the term “ground truth” refers to the accuracy of a training set's classification for supervised learning techniques. Various techniques may be used to train the models including backpropagation, statistical learning, supervised learning, semi-supervised learning, stochastic learning, or other known techniques.

(135) The ML model(s) of the shortlister component **410** may be trained and operated according to

various machine learning techniques. Such techniques may include, for example, neural networks (such as deep neural networks and/or recurrent neural networks), inference engines, trained classifiers, etc. Examples of trained classifiers include Support Vector Machines (SVMs), neural networks, decision trees, AdaBoost (short for “Adaptive Boosting”) combined with decision trees, and random forests. Focusing on SVM as an example, SVM is a supervised learning model with associated learning algorithms that analyze data and recognize patterns in the data, and which are commonly used for classification and regression analysis. Given a set of training examples, each marked as belonging to one of two categories, an SVM training algorithm builds a model that assigns new examples into one category or the other, making it a non-probabilistic binary linear classifier. More complex SVM models may be built with the training set identifying more than two categories, with the SVM determining which category is most similar to input data. An SVM model may be mapped so that the examples of the separate categories are divided by clear gaps. New examples are then mapped into that same space and predicted to belong to a category based on which side of the gaps they fall on. Classifiers may issue a “score” indicating which category the data most closely matches. The score may provide an indication of how closely the data matches the category.

(136) If the shortlister component **410** determines a natural language input is associated with multiple skill components **190/390**, only the recognizers **463** associated with those skill components **190/390** may process with respect to the natural language input. The selected recognizers **463** may process in parallel, in series, partially in parallel, etc. For example, if the shortlister component **410** determines a natural language input may relate to both a communications skill component and a music skill component, a recognizer **463** associated with the communications skill component may process in parallel, or partially in parallel, with a recognizer **463** associated with the music skill component processing.

(137) The shortlister component **410** may make binary determinations (e.g., yes or no) regarding which skill component(s) **190/390** corresponds to a natural language input. The shortlister component **410** may make such determinations using the one or more ML models described herein above. If the shortlister component **410** implements a single ML model for each skill component **190/390**, the shortlister component **410** may simply run the ML models that are associated with enabled skill components **190/390** as indicated in a profile (e.g., stored in the profile storage **170/370**) associated with the device **110** and/or user **105** that originated the natural language input.

(138) In some embodiments, the shortlister component **410** may use the widget context data **114** to determine the skill components **190/390** that correspond to the user input. The shortlister component **410** may process the widget context data **114** using one or more of the ML models described above.

(139) The shortlister component **410** may generate shortlisted skill component data **415** representing one or more skill components **190/390** that may execute in response to the user input. The skill components **190/390** included in the shortlisted skill component data **415** may be based on the widget context data **114**. The number of skill components **190/390** represented in the shortlisted skill component data **415** is configurable. In an example, the shortlisted skill component data **415** may indicate every skill component **190/390** of (or otherwise in communication with) the system **120/device 110** as well as contain an indication, for each skill component **190/390**, representing whether the skill component **190/390** is likely capable of processing in response to the natural language input. In another example, instead of indicating every skill component **190/390**, the shortlisted skill component data **415** may only indicate the skill components **190/390** that are likely capable of processing in response to the natural language input. In yet another example, the shortlister component **410** may implement thresholding such that the shortlisted skill component data **415** may indicate no more than a maximum number of skill components **190/390** that may process in response to the natural language input.

(140) In at least some embodiments, the shortlister component **410** may generate a score

representing how likely a skill component **190/390** is likely to process in response to a user input. In such embodiments, the shortlisted skill component data **415** may only include identifiers of skill components **190/390** associated with scores meeting or exceeding a threshold score.

(141) In the situation where the ASR component **150/350** outputs ASR data **402** including more than one interpretation of a spoken input, the shortlister component **410** may output different shortlisted skill component data **415** for each interpretation. Alternatively, the shortlister component **410** may output a single shortlisted skill component data **415** representing the skill components **190/390** corresponding to the different interpretations.

(142) As indicated above, the shortlister component **410** may implement thresholding such that the shortlisted skill component data **415** may indicate no more than a threshold number of skill components **190/390** (e.g., may include no more than a threshold number of skill component identifiers). If the ASR component **150/350** outputs ASR data **402** including more than one interpretation of a natural language input, the shortlisted skill component data **415** may indicate no more than a threshold number of skill components **190/390** irrespective of the number of interpretations output by the ASR component **150/350**. Alternatively or in addition, the shortlisted skill component data **415** may indicate no more than a threshold number of skill components **190/390** for each interpretation (e.g., indicating no more than five skill components **190/390** for a first interpretation, no more than five skill components **190/390** for a second interpretation, etc.).

(143) In addition to making a binary determination regarding whether a skill component **190/390** may process in response to a natural language input, the shortlister component **410** may generate scores representing likelihoods that skill components **190/390** may process in response to the natural language input. If the shortlister component **410** implements a different ML model for each skill component **190/390**, the shortlister component **410** may generate a different confidence score for each skill component-specific ML model that is run by the shortlister component **410**. For example, if the shortlister component **410** runs the ML models of every skill component **190/390** of (or otherwise in communication with) the system **120/device 110**, the shortlister component **410** may generate a respective confidence score for each skill component **190/390**. For further example, if the shortlister component **410** only runs ML models specific to skill components **190/390** that are indicated as enabled in a profile associated with the device **110** and/or user **105** (as stored in the profile storage **170/370**), the shortlister component **410** may only generate a respective confidence score for each enabled skill component **190/390**. For further example, if the shortlister component **410** implements a single ML model with skill component-specific portions, the shortlister component **410** generate a respective confidence score for each skill component **190/390** who's specifically trained portion is run. The shortlister component **410** may perform matrix vector modification to obtain confidence scores for skill components **190/390**.

(144) An example of shortlisted skill component data **415** including confidence scores may be represented as: Search skill component, 0.67 Recipe skill component, 0.62 Information skill component, 0.57

As indicated, the confidence scores output by the shortlister component **410** may be numeric values. The confidence scores output by the shortlister component **410** may alternatively be binned values (e.g., high, medium, low).

(145) The shortlister component **410** may consider other data **420** when determining which skill components **190/390** may process in response to a natural language input. The other data **420** may be character embedded prior to being input to the shortlister component **410**. The other data **420** may alternatively be embedded using other techniques.

(146) The other data **420** may include usage history data associated with the device **110** and/or user **105**. For example, a confidence score of a skill component **190/390** may be increased if natural language inputs captured by the device **110** and/or originating from the user **105** routinely relate to the skill component **190/390**. Conversely, a confidence score of a skill component **190/390** may be decreased if natural language inputs captured by the device **110** and/or originated from the user **105**

rarely relate to the skill component **190/390**.

(147) The other data **420** may indicate the skill components **190/390** that are enabled with respect to the device **110** and/or user **105** (e.g., as represented in the profile storage **170/370**). The shortlister component **410** may use such data to determine which skill component-specific ML models to run. That is, the shortlister component **410** may determine to only run the ML models associated with enabled skill components **190/390**. The shortlister component **410** may alternatively use such data to alter skill component confidence scores represented in the shortlisted skill component data **415**. As an example, considering two skill components **190/390**, one enabled and another unenabled, the shortlister component **410** may run a first ML model (or ML model portion) specific to the unenabled skill component as well as a second ML model (or ML model portion) specific to the enabled skill component. The shortlister component **410** may initially determine a confidence score of 0.60 for each of the unenabled skill component and the enabled skill component. The shortlister component **410** may then alter those confidence scores based on which skill component is enabled. For example, the shortlister component **410** may increase the confidence score associated with the enabled skill component and/or decrease the confidence score associated with the unenabled skill component.

(148) A user **105** may provide the system **120** with indications of which skill components **190/390** are enabled (e.g., authorized to execute using data associated with the user **105**). Such indications may be stored in the profile storage **170/370**. The shortlister component **410** may determine whether profile data associated with the user **105** and/or device **110** that originated the natural language input includes indications of enabled skill components **190/390**.

(149) The other data **420** may indicate a type of the device **110**. The type of the device **110** may indicate the input/output capabilities of the device **110**. For example, the device **110** may include a display, may be headless (e.g., displayless), may be mobile or stationary, may include audio playback capabilities, may include a camera, etc. The shortlister component **410** may use such other data **420** to determine which skill component-specific ML models (or ML model portions) to run. For example, if the device **110** is displayless, the shortlister component **410** may determine not to run ML models (or ML model portions) specific to skill components **190/390** that output video data. The shortlister component **410** may alternatively use such other data **420** to alter skill component confidence scores represented in the shortlisted skill component data **415**. As an example, considering two skill components **190/390**, one that outputs audio data and another that outputs video data, the shortlister component **410** may run a first ML model (or portion of a ML model) specific to the skill component that generates audio data as well as a second ML model (or portion of a ML model) specific to the skill component that generates video data. The shortlister component **410** may initially determine a confidence score of 0.60 for each of the skill components. The shortlister component **410** may then alter the original confidence scores based on the type of the device **110**. For example, if the device **110** is displayless, the shortlister component **410** may increase the confidence score associated with the skill component that generates audio data and/or decrease the confidence score associated with the skill component that generates video data.

(150) The type of device data represented in the other data **420** may represent output capabilities of the device **110** to be used to output content to the user **105**, which may not necessarily be the device **110** that captured the natural language input. For example, a displayless device **110** may receive a natural language input corresponding to “play Game of Thrones”. The system **120/device 110** may determine a smart TV, or other device including or otherwise associated with a display, is to be used to output multimedia content (e.g., audio and video) corresponding to the title “Game of Thrones”.

(151) The other data **420** may include data indicating a speed of the device **110**, a location of the device **110**, or other mobility data. For example, the device **110** may correspond to a vehicle including a display. If the vehicle is moving, the shortlister component **410** may decrease the confidence score associated with a skill component **190/390** that generates video data as it may be

undesirable to output video content to the user **105** while in a moving vehicle. The vehicle may output data to the system **120** indicating when the vehicle is in motion.

(152) The other data **420** may include data indicating a currently invoked skill component **190/390** (e.g., a skill component **190/390** that was processing to output content to the user **105** when the device **110** received the natural language input). For example, the user **105** may speak a first (e.g., a previous) natural language input causing the system **120** to invoke a music skill component to output music to the user **105**. As the music is being output to the user **105**, the system **120** may receive a second (e.g., a current) natural language input. The shortlister component **410** may use such other data **420** to alter skill component confidence scores represented in the shortlisted skill component data **415**. For example, the shortlister component **410** may run a first ML model (or ML model portion) specific to a first skill component as well as a second ML model (or ML model portion) specific to a second skill component. The shortlister component **410** may initially determine a confidence score of 0.60 for each of the first and second skill components. The shortlister component **410** may then alter the original confidence scores based on the first skill component being invoked to output content while the current natural language input was received. Based on the first skill component being invoked, the shortlister component **410** may increase the confidence score associated with the first skill component and/or decrease the confidence score associated with the second skill component.

(153) The thresholding implemented with respect to the shortlisted skill component data **415** and the different types of other data **420** considered by the shortlister component **410** are configurable. For example, the shortlister component **410** may update confidence scores as more other data **420** is considered.

(154) The shortlister component **410** may cause the NLU component **160/360** to invoke only a subset of the recognizers **463** associated with skill components **190/390** represented in the shortlisted skill component **415** as being likely to process in response to the natural language input. If the shortlister component **410** generates the shortlisted skill component data **415** to include confidence scores, the shortlister component **410** may cause the NLU component **160/360** to invoke only recognizers **463** associated with skill components **190/390** associated with confidence scores satisfying (e.g., meeting or exceeding) a condition (e.g., a threshold confidence score).

(155) The NLU component **160/360** may compile NLU hypotheses, output by multiple recognizers **463**, into cross-recognizer hypothesis data **440** (illustrated in FIG. 4). Each NLU hypothesis may be associated with a respective score indicating a likelihood that the NLU hypothesis corresponds to the domain, one or more skill components **190/390**, etc. associated with the recognizer **463** from which the NLU hypothesis was output. For example, the cross-recognizer hypothesis data **440** may be represented as, with each line corresponding to a different NLU hypothesis: [0.95] Intent: <PlayMusic> AlbumName: GameOfThrones [0.70] Intent: <PlayVideo> VideoTitle: GameOfThrones.

While the foregoing illustrates cross-recognizer hypothesis data **440** include two NLU hypotheses, it will be appreciated that the cross-recognizer hypothesis data **440** may include differing numbers of NLU hypotheses with departing from the present disclosure.

(156) The NLU component **160/360** may send the cross-recognizer hypothesis data **440** to a pruning component **450**, which sorts the NLU hypotheses, represented in the cross-recognizer hypothesis data **440**, according to their respective scores. The pruning component **450** may then perform score thresholding with respect to the cross-recognizer hypothesis data **440**. For example, the pruning component **450** may select NLU hypotheses, represented in the cross-recognizer hypothesis data **440**, associated with scores satisfying (e.g., meeting and/or exceeding) a condition (e.g., a threshold score). The pruning component **450** may additionally or alternatively perform number of NLU hypothesis thresholding. For example, the pruning component **450** may select a threshold number of top-scoring NLU hypotheses represented in the cross-recognizer hypothesis data **440**.

(157) The pruning component **450** may generate cross-recognizer hypothesis data **460** including the selected NLU hypotheses. The purpose of the pruning component **450** is to create a reduced list of NLU hypotheses so that downstream, more resource intensive, processes may only operate on NLU hypotheses that most likely represent the natural language input.

(158) The NLU component **160/360** may include a light slot filler component **452** that takes text from slots, represented in the NLU hypotheses output by the pruning component **450**, and alters it to make the text more easily processed by downstream components. The light slot filler component **452** may perform low latency operations that do not involve heavy operations such as reference to a knowledge base. The purpose of the light slot filler component **452** is to replace words with other words or values that may be more easily understood by downstream components. For example, if a NLU hypothesis includes the word “tomorrow,” the light slot filler component **452** may replace the word “tomorrow” with an actual date for purposes of downstream processing. Similarly, the light slot filler component **452** may replace the word “CD” with “album” or “compact disc.” The replaced text is then included in the cross-recognizer hypothesis data **460**.

(159) The cross-recognizer hypothesis data **460** may be sent to an entity resolution component **470**. The entity resolution component **470** can apply rules or other instructions to standardize labels or tokens in the NLU hypotheses represented in the cross-recognizer hypothesis data **460**. The precise transformation may depend on the skill component **190/390**, domain, etc. to which a NLU hypothesis relates. For example, for a travel skill component-specific NLU hypothesis, the entity resolution component **470** may transform text corresponding to “Boston airport” to the standard “BOS” three-letter code referring to the airport. The entity resolution component **470** can refer to a knowledge base that is used to specifically identify the precise entity referred to in each slot of each NLU hypothesis represented in the cross-recognizer hypothesis data **460**.

(160) Specific intent/slot combinations may also be tied to a particular source, which may then be used to resolve the text. Referring to the example natural language input “play songs by the stones,” the entity resolution component **470** may reference a personal music catalog, a user profile, or the like (for example stored in the profile storage **170/370**). The entity resolution component **470** may output NLU hypotheses, altered from the cross-recognizer hypothesis data **460**, that include more detailed information (e.g., entity identifiers) about the specific entities mentioned in the slots and/or more detailed slot data that can eventually be used by a skill component **190/390** to perform an action responsive to the natural language input. The NLU component **160/360** may include multiple entity resolution components **470** that are each specific to one or more different skill components **190/390**, domains, etc.

(161) In some embodiments, the entity resolution component **470** may use the widget context data **114** to determine the specific entity corresponding to a slot/entity represented in the user input.

(162) The NLU component **160/360** may include a ranker component **490** that assigns a particular score to each NLU hypothesis output by the entity resolution component **470**. The ranker component **490** may alter the score of a particular NLU hypothesis based on whether the NLU hypothesis has unfilled slots. For example, if a first NLU hypothesis includes slots that are all filled/resolved, the ranker component **490** may assign the first NLU hypothesis a higher score than a second NLU hypothesis including at least one slot that is unfilled/unresolved by the entity resolution component **470**.

(163) The ranker component **490** may apply re-scoring, biasing, or other techniques. To do so, the ranker component **490** may consider not only the data output by the entity resolution component **470**, but may also other data **491**. The other data **491** may include a variety of data.

(164) For example, the other data **491** indicate skill component **190/390** rating or popularity. For example, if a skill component **190/390** has a high rating, the ranker component **490** may increase the score of a NLU hypothesis associated with that skill component **190/390**, and vice versa.

(165) The other data **491** may indicate skill components **190/390** that have been enabled by the user **105**. For example, the ranker component **490** may assign higher scores to NLU hypotheses

associated with enabled skill components than NLU hypotheses associated with skill components that have not been enabled by the user **105**.

(166) The other data **491** include system usage history (e.g., specific to the user **105**), such as if the user **105** regularly invokes a particular skill component **190/390** or does so at particular times of day. The other data **491** may indicate a present date, a present time, location of the device **110**, weather data, a type of the device **110**, a user identifier of the user **105**, context data, as well as other data. For example, the ranker component **490** may consider when any particular skill component **190/390** is currently active with respect to the present user **105** and/or device **110** (e.g., music being output by the skill component **190/390** when the current natural language input is received).

(167) The ranker component **490** may output NLU output data **485** including one or more NLU hypotheses. The NLU component **160/360** may send the NLU output data **485** to the orchestrator component **130**.

(168) Referring to FIG. 5, details of how the skill selection component **165/365** may be configured are described. When the skill selection component **165/365** receives NLU output data **485** from the NLU component **160**, the NLU output data **485** may be sent to a skill pre-response component **510**. The skill pre-response component **510** is configured to query skills, associated with the NLU output data **485** as to whether the skills are able to process the NLU hypothesis with which they are associated. In some embodiments, the skill pre-response component **510** may query skills associated with the multi-modal widgets **112** as to whether the skills are able to process an NLU hypothesis.

(169) The skill pre-response component **510** sends a pre-response query to each skill component **190/390** associated with a NLU hypothesis represented in the NLU output data **485**. A pre-response query may include a NLU hypothesis (associated with the skill component **190/390** receiving the pre-response query), portion of the widget context data **114** that corresponds to the queried skill, and optionally other context data corresponding to the user input.

(170) A skill component **190/390** may determine, based on a received pre-response query (and optionally other data available to the skill component **190/390**), whether the skill **190/390** could respond to the natural language input. For example, a skill component **190/390** may generate a pre-response indicating the skill component **190/390** can respond to the natural language input, indicating the skill component **190/390** may be able to respond to the natural language input (e.g., the indicating the skill component **190/390** needs more data to determine whether the skill component **190/390** can respond to the natural language input), or indicating the skill component **190/390** cannot respond to the natural language input (e.g., due to present processing load).

(171) In situations where a skill's pre-response indicates the skill component **190/390** can or may be able to respond to the natural language input, the pre-response may also include various other data representing a strength of the skill's potential response to the natural language input. Such other data may positively influence the skill's ranking as described herein below. For example, such other data may indicate capabilities (e.g., output capabilities or components such as a connected screen, loudspeaker, etc.) of a device to be used to output the skill's response; pricing data corresponding to a product or service the natural language input is requesting be purchased (or is requesting purchase data for); availability of a product the natural language input is requesting be purchased; whether there are shipping fees for a product the natural language input is requesting be purchased; whether the user **105** already has a profile and/or subscription with the skill component **190/390**; that the user **105** does not have a subscription with the skill component **190/390**, but that there is a free trial/tier the skill component **190/390** is offering; with respect to a taxi skill, a cost of a trip based on start and end locations, how long the user **105** would have to wait to be picked up, etc.; and/or other data available to the skill component **190/390** that is related to the skill's processing of the natural language input. In at least some embodiments, a skill's pre-response may include an indicator (e.g., a flag) representing a strength of the skill's ability to personalize its

response to the user **105**.

(172) In at least some embodiments, a skill's pre-response may be configured to a schema pre-defined by the system **120**. By the system **120** requiring pre-responses to conform to a specific schema (e.g., by requiring skill components **190/390** to only be able to provide certain types of data in pre-responses), the system **120** may onboard new skill components **190/390** into the skill selection functionality described herein without needing to reconfigure the skill selection component **165/365** each time a new skill component **190/390** is onboarded. Moreover, requiring pre-responses to conform to a schema limits the amount of values needed to be used to train and implement a machine learning model for ranking skill components **190/390**, as described herein below.

(173) The skill pre-response component **510** may receive pre-response data from the queried skills, and output data **515** including the skill identifiers associated with respective pre-responses. As indicated above, a skill's pre-response data may indicate the skill component **190/390** would be unable to process if called to execute with respect to the natural language input. In at least some embodiments, the skill pre-response component **510** may only include, in the data **515**, the skill identifiers associated with pre-response data that does not indicate an associated skill component **190/390** would be unable to process. The data **515**, output by the skill pre-response component **510**, may not include any confidence scores representing associated skill component **190/390** should be used to respond to the natural language input.

(174) The data **515**, output by the skill pre-response component **510**, may be input to a skill ranker **520** of the skill selection component **165/365**. The skill ranker **520** may be configured to rank skill components **190/390** using a machine learning model. In at least some embodiments, the machine learning model may be a deep neural network (DNN).

(175) The skill ranker **520** also receives context data corresponding to the natural language input. Such context data may indicate, for example, one or more preferences (e.g., skill **125** preferences) of the user **105** (as represented in a user profile or group profile stored in the profile storage **170**), one or more domains corresponding to the natural language input, one or more intents representing the natural language input, NER entity values as represented in the NLU output data, a measure of success of entity resolution performed with respect to the natural language input, text representing the natural language input, rankings of individual NLU hypotheses within the NLU output data, a device type of the device **110** that received the natural language input, whether the device **110** was outputting audio when the natural language input was received, whether the device **110** was outputting video when the natural language input was received, and/or other context data available to the skill ranker **520**. The skill ranker **520** may also receive the widget context data **114**.

(176) Using the ML model, the skill ranker **520** ranks the skill component **190/390** based on the data **515**, the widget context data **114**, and the context data. Things that may increase a skill's ranking include, for example, that the skill component **190/390** is associated with pre-response data indicating the skill component **190/390** can generate a response (to the natural language input) that is personalized to the user **105**, that a NLU hypothesis corresponding to the skill component **190/390** is associated with a NLU confidence score satisfying (e.g., meeting or exceeding) a condition (e.g., a threshold NLU confidence score), that the skill component **190/390** corresponds to a music skill and the device **110** was outputting music when the natural language input was received, that the skill component **190/390** corresponds to a video skill and the device **110** was outputting video when the natural language input was received, etc. Things that may decrease a skill's ranking include, for example, that the skill component **190/390** is associated with pre-response data indicating the skill component **190/390** cannot generate a response (to the natural language input) that is personalized to the user **105**, that a NLU hypothesis corresponding to the skill component **190/390** is associated with a NLU confidence score failing to satisfy (e.g., falling below) a condition (e.g., a threshold NLU confidence score), that the skill component **190/390** corresponds to a video skill and the device **110** was outputting music when the natural language

input was received, that the skill component **190/390** corresponds to a music skill and the device **110** was outputting video when the natural language input was received, etc. Other example are possible and will depend on the context data available to the skill ranker **520**.

(177) The skill ranker **520** outputs data **525** including the skill identifiers associated with respective confidence scores, where a confidence score represents a strength with which the skill ranker **520** recommends the associated skill component **190/390** is to execute to respond to the natural language input. Such a confidence score may be a numeric score (e.g., between 0 and 1) or a binned score (e.g., low, medium, high). The data **525** may also include one or more widget identifiers (and/or widget instance identifiers) corresponding to the multi-modal widgets **112**, when the skill component **190/390** included in the data **525** corresponds to the multi-modal widgets **112** active or enabled at the device **110**. As such, the data **525** may include skill identifiers associated with one or more skill components **190/390** that has a multi-modal widget active or enabled at the device **110**, and/or may include skill identifiers associated with one or more skill components **190/390** that do not have a multi-modal widget active or enabled at the device **110**.

(178) A decider component **530** may receive the data **525** output by the skill ranker **520**. The decider component **530** may be configured to generate an execution plan **555** based on the data **525** (representing the skill identifiers associated with respective confidence scores) and context data (e.g., representing NLU hypotheses associated with the skill component identifiers and the widget context data **114**).

(179) In some embodiments, the decider component **530** may select the skill component **190/390** associated with the highest confidence score in the data **525**. In some embodiments, the decider component **530** may select the skill component **190/390** associated with a confidence score that satisfies a condition (e.g., exceeds a threshold value). In some embodiments, the decider component **530** may select the skill component **190/390** associated with a confidence score that satisfies a threshold condition and satisfies a condition corresponding to context data, such as the widget context data **114**.

(180) The decider component **530** may generate an execution plan **555** including the skill identifier, of the skill component **190/390** selected by the decider component **530** to execute in response to the natural language input, the NLU hypothesis associated with the skill identifier, a portion of the widget context data **114** corresponding to the selected skill component **190/390** (i.e. the widget context data **116**) and a command to execute/invoke the skill component **190/390** corresponding to the skill identifier. The decider component **530** may send the execution plan **555** to the skill request router component **188**, which may cause the skill component **190/390** (corresponding to the skill identifier in the execution plan **555**) to process the NLU hypothesis to determine output data responsive to the natural language input.

(181) FIG. **6** is a block diagram conceptually illustrating a device **110** of the system **100**. FIG. **7** is a block diagram conceptually illustrating example components of a remote device, such as the system **120** and the skill(s) system **125**. A system (**120/125**) may include one or more servers. A “server” as used herein may refer to a traditional server as understood in a server/client computing structure but may also refer to a number of different computing components that may assist with the operations discussed herein. For example, a server may include one or more physical computing components (such as a rack server) that are connected to other devices/components either physically and/or over a network and is capable of performing computing operations. A server may also include one or more virtual machines that emulates a computer system and is run on one or across multiple devices. A server may also include other combinations of hardware, software, firmware, or the like to perform operations discussed herein. The system (**120/125**) may be configured to operate using one or more of a client-server model, a computer bureau model, grid computing techniques, fog computing techniques, mainframe techniques, utility computing techniques, a peer-to-peer model, sandbox techniques, or other computing techniques.

(182) Multiple systems (**120/125**) may be included in the system **100** of the present disclosure,

such as, one or more systems **120** and/or one or more skills **125**. In operation, each of these systems may include computer-readable and computer-executable instructions that reside on the respective device (**120/125**), as will be discussed further below.

(183) Each of these devices (**110/120/125**) may include one or more controllers/processors (**604/704**), which may each include a central processing unit (CPU) for processing data and computer-readable instructions, and a memory (**606/706**) for storing data and instructions of the respective device. The memories (**606/706**) may individually include volatile random access memory (RAM), non-volatile read only memory (ROM), non-volatile magnetoresistive memory (MRAM), and/or other types of memory. Each device (**110/120/125**) may also include a data storage component (**608/708**) for storing data and controller/processor-executable instructions. Each data storage component (**608/708**) may individually include one or more non-volatile storage types such as magnetic storage, optical storage, solid-state storage, etc. Each device (**110/120/125**) may also be connected to removable or external non-volatile memory and/or storage (such as a removable memory card, memory key drive, networked storage, etc.) through respective input/output device interfaces (**602/702**).

(184) Computer instructions for operating each device (**110/120/125**) and its various components may be executed by the respective device's controller(s)/processor(s) (**604/704**), using the memory (**606/706**) as temporary “working” storage at runtime. A device's computer instructions may be stored in a non-transitory manner in non-volatile memory (**606/706**), storage (**608/708**), or an external device(s). Alternatively, some or all of the executable instructions may be embedded in hardware or firmware on the respective device in addition to or instead of software.

(185) Each device (**110/120/125**) includes input/output device interfaces (**602/702**). A variety of components may be connected through the input/output device interfaces (**602/702**), as will be discussed further below. Additionally, each device (**110/120/125**) may include an address/data bus (**624/724**) for conveying data among components of the respective device. Each component within a device (**110/120/125**) may also be directly connected to other components in addition to (or instead of) being connected to other components across the bus (**624/724**).

(186) Referring to FIG. 6, the device **110** may include input/output device interfaces **602** that connect to a variety of components such as an audio output component such as a speaker **612**, a wired headset or a wireless headset (not illustrated), or other component capable of outputting audio. The device **110** may also include an audio capture component. The audio capture component may be, for example, a microphone **620** or array of microphones, a wired headset or a wireless headset (not illustrated), etc. If an array of microphones is included, approximate distance to a sound's point of origin may be determined by acoustic localization based on time and amplitude differences between sounds captured by different microphones of the array. The device **110** may additionally include a display **616** for displaying content. The device **110** may further include a camera **618**.

(187) Via antenna(s) **614**, the input/output device interfaces **602** may connect to one or more networks **199** via a wireless local area network (WLAN) (such as Wi-Fi) radio, Bluetooth, and/or wireless network radio, such as a radio capable of communication with a wireless communication network such as a Long Term Evolution (LTE) network, WiMAX network, 3G network, 4G network, 5G network, etc. A wired connection such as Ethernet may also be supported. Through the network(s) **199**, the system may be distributed across a networked environment. The I/O device interface (**602/702**) may also include communication components that allow data to be exchanged between devices such as different physical servers in a collection of servers or other components.

(188) The components of the device **110**, the system **120**, and/or skill **125** may include their own dedicated processors, memory, and/or storage. Alternatively, one or more of the components of the device **110**, the system **120**, and/or skill **125** may utilize the I/O interfaces (**602/702**), processor(s) (**604/704**), memory (**606/706**), and/or storage (**608/708**) of the device **110**, the system **120**, and/or skill **125**, respectively.

(189) As noted above, multiple devices may be employed in a single system. In such a multi-device system, each of the devices may include different components for performing different aspects of the system's processing. The multiple devices may include overlapping components. The components of the device **110**, the system **120**, and the skill **125**, as described herein, are illustrative, and may be located as a stand-alone device or may be included, in whole or in part, as a component of a larger device or system.

(190) As illustrated in FIG. **8**, multiple devices (**110a-110j**, **120**, **125**) may contain components of the system and the devices may be connected over a network(s) **199**. The network(s) **199** may include a local or private network or may include a wide network such as the Internet. Devices may be connected to the network(s) **199** through either wired or wireless connections. For example, a speech-controllable device **110a**, a smart phone **110b**, a smart watch **110c**, a tablet computer **110d**, a vehicle **110e**, a speech-controllable display device **110f**, a smart television **110g**, a washer/dryer **110h**, a refrigerator **110i**, and/or a microwave **110j** may be connected to the network(s) **199** through a wireless service provider, over a Wi-Fi or cellular network connection, or the like. Other devices are included as network-connected support devices, such as the system **120**, the skill **125**, and/or others. The support devices may connect to the network(s) **199** through a wired connection or wireless connection.

(191) The concepts disclosed herein may be applied within a number of different devices and computer systems, including, for example, general-purpose computing systems, speech processing systems, and distributed computing environments.

(192) The above aspects of the present disclosure are meant to be illustrative. They were chosen to explain the principles and application of the disclosure and are not intended to be exhaustive or to limit the disclosure. Many modifications and variations of the disclosed aspects may be apparent to those of skill in the art. Persons having ordinary skill in the field of computers and speech processing should recognize that components and process steps described herein may be interchangeable with other components or steps, or combinations of components or steps, and still achieve the benefits and advantages of the present disclosure. Moreover, it should be apparent to one skilled in the art, that the disclosure may be practiced without some or all of the specific details and steps disclosed herein.

(193) Aspects of the disclosed system may be implemented as a computer method or as an article of manufacture such as a memory device or non-transitory computer readable storage medium. The computer readable storage medium may be readable by a computer and may comprise instructions for causing a computer or other device to perform processes described in the present disclosure. The computer readable storage medium may be implemented by a volatile computer memory, non-volatile computer memory, hard drive, solid-state memory, flash drive, removable disk, and/or other media. In addition, components of system may be implemented as in firmware or hardware, such as an acoustic front end (AFE), which comprises, among other things, analog and/or digital filters (e.g., filters configured as firmware to a digital signal processor (DSP)).

(194) Conditional language used herein, such as, among others, “can,” “could,” “might,” “may,” “e.g.,” and the like, unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps. Thus, such conditional language is not generally intended to imply that features, elements, and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without other input or prompting, whether these features, elements, and/or steps are included or are to be performed in any particular embodiment. The terms “comprising,” “including,” “having,” and the like are synonymous and are used inclusively, in an open-ended fashion, and do not exclude additional elements, features, acts, operations, and so forth. Also, the term “or” is used in its inclusive sense (and not in its exclusive sense) so that when used, for example, to connect a list of elements, the term “or” means one, some, or all of the elements in the

list.

(195) Disjunctive language such as the phrase “at least one of X, Y, Z,” unless specifically stated otherwise, is understood with the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present.

(196) As used in this disclosure, the term “a” or “one” may include one or more items unless specifically stated otherwise. Further, the phrase “based on” is intended to mean “based at least in part on” unless specifically stated otherwise.

Claims

1. A computer-implemented method, comprising: receiving first input data corresponding to a first natural language input; receiving first data corresponding to a first graphical element being displayed on a screen of a device, the first graphical element corresponding to first content capable of being interacted with through at least one of physical input or natural language input; receiving second data corresponding to second graphical element being displayed on the screen of the device, the second graphical element corresponding to second content capable of being interacted with through at least one of physical input or natural language input; determining, using the first input data, the first data and the second data, that the first natural language input corresponds to the first graphical element; and based at least in part on the first natural language input corresponding to the first graphical element, causing an action to be performed with regard to the first content.
2. The computer-implemented method of claim 1, further comprising: receiving first application data from a first application; and causing the first application data to be displayed on the screen as part of the first graphical element.
3. The computer-implemented method of claim 1, further comprising: receiving second data representing speech of the first natural language input; and performing speech processing using the second data to determine the first input data.
4. The computer-implemented method of claim 1, wherein determining that the first natural language input corresponds the first content comprises performing natural language processing.
5. The computer-implemented method of claim 4, wherein performing natural language processing comprises processing the first data and the second data.
6. The computer-implemented method of claim 1, further comprising: determining the first natural language input mentions an entity; and determining that the entity is represented in the first data, wherein determining that the first natural language input corresponds the first content is based at least in part on the entity being represented in the first content.
7. The computer-implemented method of claim 1, wherein the first data comprises time data corresponding to a time of a most recent user input corresponding to the first graphical element.
8. The computer-implemented method of claim 1, wherein the first data comprises time data corresponding to a time of update of the first graphical element.
9. The computer-implemented method of claim 1, wherein the first data comprises position information corresponding to a position of the first graphical element with respect to the screen.
10. The computer-implemented method of claim 1, wherein the first data comprises an identifier of a source of the first content and wherein performing the action comprises sending an instruction to the source.
11. A system comprising: at least one processor; and at least one memory comprising instructions that, when executed by the at least one processor, cause the system to: receive first input data corresponding to a first natural language input; receive first data corresponding to a first graphical element being displayed on a screen of a device, the first graphical element corresponding to first content capable of being interacted with through at least one of physical input or natural language

input; receive second data corresponding to second graphical element being displayed on the screen of the device, the second graphical element corresponding to second content capable of being interacted with through at least one of physical input or natural language input; determine, using the first input data, the first data and the second data, that the first natural language input corresponds to the first graphical element; and based at least in part on the first natural language input corresponding to the first graphical element, cause an action to be performed with regard to the first content.

12. The system of claim 11, wherein the at least one memory further comprises instructions that, when executed by the at least one processor, further cause the system to: receive first application data from a first application; and cause the first application data to be displayed on the screen as part of the first graphical element.

13. The system of claim 11, wherein the at least one memory further comprises instructions that, when executed by the at least one processor, further cause the system to: receive second data representing speech of the first natural language input; and perform speech processing using the second data to determine the first input data.

14. The system of claim 11, wherein the instructions that cause the system to determine that the first natural language input corresponds the first content comprise instructions that, when executed by the at least one processor, cause the system to perform natural language processing.

15. The system of claim 14, wherein the instructions that cause the system to determine perform natural language processing comprise instructions that, when executed by the at least one processor, cause the system to process the first data and the second data.

16. The system of claim 11, wherein the at least one memory further comprises instructions that, when executed by the at least one processor, further cause the system to: determine the first natural language input mentions an entity; and determine that the entity is represented in the first data, wherein determining that the first natural language input corresponds the first content is based at least in part on the entity being represented in the first content.

17. The system of claim 11, wherein the first data comprises time data corresponding to a time of a most recent user input corresponding to the first graphical element.

18. The system of claim 11, wherein the first data comprises time data corresponding to a time of update of the first graphical element.

19. The system of claim 11, wherein the first data comprises position information corresponding to a position of the first graphical element with respect to the screen.

20. The system of claim 11, wherein the first data comprises an identifier of a source of the first content and wherein the instructions that cause the system to perform the action comprise instructions that, when executed by the at least one processor, cause the system to send an instruction to the source.
