

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386654
Kind Code	B1
Date of Patent	August 12, 2025
Inventor(s)	Govindaraju; Venkatraman et al.

Crash resilience in shared memory processing systems

Abstract

Improved crash resilience is implemented in shared memory processing systems. A processing application including a plurality processing groups may be configured, and a processing group including a plurality of processes may allocate memory from a global shared memory to a group-private shared memory arena. Then, an abnormal execution condition for a failing process of the processing group may be detected during execution of a processing job. Responsive to the detection, a recovery of execution of the processing group may be performed while continuing execution of processing groups of the application. The recovery may include restoring the processing group to a normal execution state and returning the allocated memory of the group-private memory arena to the global shared memory.

Inventors: Govindaraju; Venkatraman (Fremont, CA), Sokolov; Pavel (Menlo Park, CA), Chainani; Naresh (Mountain View, CA)

Applicant: Amazon Technologies, Inc. (Seattle, WA)

Family ID: 96661891

Assignee: Amazon Technologies, Inc. (Seattle, WA)

Appl. No.: 17/809873

Filed: June 29, 2022

Publication Classification

Int. Cl.: G06F9/48 (20060101); G06F9/50 (20060101); G06F9/54 (20060101)

U.S. Cl.:

CPC G06F9/485 (20130101); G06F9/505 (20130101); G06F9/5072 (20130101); G06F9/5077 (20130101); G06F9/544 (20130101)

Field of Classification Search

CPC: G06F (9/485)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7219149	12/2006	Ofir	709/224	H04M 11/06
9218246	12/2014	Nagesharao et al.	N/A	N/A
10542068	12/2019	Vermeulen et al.	N/A	N/A
10866869	12/2019	Li	N/A	N/A
11030055	12/2020	Gupta et al.	N/A	N/A
2015/0269111	12/2014	Bak	711/147	G06F 13/1663

Primary Examiner: Sun; Charlie

Attorney, Agent or Firm: Kowert, Hood, Munyon, Rankin & Goetzel, P.C.

Background/Summary

BACKGROUND

(1) In shared memory data processing systems, crashes in one process or task may force all processes or tasks to be restarted because crashed process may leave the shared state in an inconsistent state. As these data processing systems grow larger, restarting becomes more computationally expensive as well as increasingly common, thus limiting the scale of such systems. To improve scalability in such systems, a means of reducing exposure to corruption of shared memory state is needed.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 illustrates a logical block diagram of a processing system implementing improved crash resilience in shared memory processing systems, according to some embodiments.
- (2) FIG. 2 is a logical block diagram illustrating an example of query processing groups performing queries to database data, according to some embodiments.
- (3) FIG. 3A is a logical block diagram of a conventional shared memory processing system, according to some embodiments.
- (4) FIG. 3B is a logical block diagram of a shared memory processing system with improved crash resilience, according to some embodiments.
- (5) FIG. 4 is a flow diagram illustrating improved crash resilience, according to some embodiments.
- (6) FIG. 5A is a logical block diagram illustrating fault signaling of a computing node in a shared memory processing system with improved crash resilience, according to some embodiments.
- (7) FIG. 5B is a logical block diagram illustrating fault signaling of a leader node in a shared memory processing system with improved crash resilience, according to some embodiments.

(8) FIG. 6 is a flow diagram illustrating fault signaling in a shared memory processing system with improved crash resilience, according to some embodiments.

(9) FIG. 7 is a flow diagram illustrating updating global memory state in a shared memory processing system with improved crash resilience, according to some embodiments.

(10) FIG. 8 illustrates an example system that implements the various methods, techniques, and systems described herein, according to some embodiments.

(11) While embodiments are described herein by way of example for several embodiments and illustrative drawings, those skilled in the art will recognize that the embodiments are not limited to the embodiments or drawings described. It should be understood that the drawings and detailed description thereto are not intended to limit embodiments to the particular form disclosed, but on the contrary, the intention is to cover all modifications, equivalents and alternatives falling within the spirit and scope as defined by the appended claims. The headings used herein are for organizational purposes only and are not meant to be used to limit the scope of the description or the claims. As used throughout this application, the word “may” is used in a permissive sense (i.e., meaning having the potential to), rather than the mandatory sense (i.e., meaning must). The words “include,” “including,” and “includes” indicate open-ended relationships and therefore mean including, but not limited to. Similarly, the words “have,” “having,” and “has” also indicate open-ended relationships, and thus mean having, but not limited to. The terms “first,” “second,” “third,” and so forth as used herein are used as labels for nouns that they precede, and do not imply any type of ordering (e.g., spatial, temporal, logical, etc.) unless such an ordering is otherwise explicitly indicated.

(12) “Based On.” As used herein, this term is used to describe one or more factors that affect a determination. This term does not foreclose additional factors that may affect a determination. That is, a determination may be solely based on those factors or based, at least in part, on those factors. Consider the phrase “determine A based on B.” While B may be a factor that affects the determination of A, such a phrase does not foreclose the determination of A from also being based on C. In other instances, A may be determined based solely on B.

(13) The scope of the present disclosure includes any feature or combination of features disclosed herein (either explicitly or implicitly), or any generalization thereof, whether or not it mitigates any or all of the problems addressed herein. Accordingly, new claims may be formulated during prosecution of this application (or an application claiming priority thereto) to any such combination of features. In particular, with reference to the appended claims, features from dependent claims may be combined with those of the independent claims and features from respective independent claims may be combined in any appropriate manner and not merely in the specific combinations enumerated in the appended claims.

DETAILED DESCRIPTION

(14) In shared memory data processing systems, crashes of one process or task may force other processes or tasks to be restarted due to an inconsistent state of global shared memory. As the data processing system grows larger, restarting becomes more computationally expensive as well as increasingly common, thus limiting the scale of such systems.

(15) Rather than restart the entire data processing system upon detection of an abnormal execution condition, the misbehaving processing job may instead be canceled or aborted and a processing group assigned to execute the affected job may be restarted, depending on the process demonstrating the abnormal execution condition. This may be performed almost entirely within a signal handler(s) of the affected process or processing group itself. In this way, processing clusters implementing such data processing system may benefit from improved scalability. However, to ensure stability and correctness of the data processing system, shared memory state that has been modified by the affected processing group must be restored. To accomplish this, modifications to shared memory processing may be employed to reduce the impact of shared memory state by processing groups that may exhibit abnormal execution conditions.

(16) Systems and methods for implementing improved crash resilience in shared memory processing systems are described herein. A processing application including a plurality of processing groups may be configured, and a processing group including a plurality of processes may allocate memory from a global shared memory to a group-private shared memory arena. Then, an abnormal execution condition for a failing process of the processing group may be detected during execution of a processing job. Responsive to the detection, a recovery of execution of the processing group may be performed while continuing execution of processing groups of the application. The recovery may include restoring the processing group to a normal execution state and returning the allocated memory of the group-private memory arena to the global shared memory.

(17) FIG. 1 illustrates a logical block diagram of a processing system implementing improved crash resilience in shared memory processing systems, according to some embodiments. A computing cluster **100** of one or more computing nodes **101** may individually include one or more processors **102** and a memory **104** to implement a data processing application **110**, in some embodiments. An example of such computing nodes is provided in the computer system embodiment illustrated in FIG. 8 and described below.

(18) The application **110** may include a process manager **140** that may manage a number of processing groups **120** that may execute processing jobs of the application **110**. These processing groups **120** may individually include a leader process **122** and optionally one or more compute processes **124**, in various embodiments.

(19) The application **110** may further include a shared memory manager **130**, in some embodiments. This shared memory may be shared across all processing groups **120** to enable the various processing groups to access and affect changes to global data of the application. For example, the application **110** may implement portions of a database service, as discussed further below in FIG. 2, with the shared memory manager **130** enabling database query processing groups to access and affect changes to various databases, in some embodiments.

(20) The shared memory manager **130** may include a memory pool **132** of shared memory pages, as shown in further detail in FIG. 3 below. In addition, shared memory manager **130** may include a number of shared memory arenas including a public arena **134** and respective private arenas **136** for the respective processing groups **120**. By isolating shared memory updates to group-private memory arenas **136** and minimizing updates to the public arena **134**, as discussed further in FIGS. 3-8 below, crash resilience in shared memory processing systems may be improved, in various embodiments.

(21) Please note that the previous description of a data processing application is a logical description and thus is not to be construed as limiting as to the implementation of a data processing systems or portions thereof.

(22) This specification continues with a general description of a provider network that implements multiple different services, including a database service and storage service, which may implement detecting idle periods for management actions at processing clusters for managed databases. Then various examples of the database service and storage service, including different components/modules, or arrangements of components/module that may be employed as part of implementing the services are discussed. A number of different methods and techniques to implement detecting idle periods for management actions at processing clusters for managed databases are then discussed, some of which are illustrated in accompanying flowcharts. Finally, a description of an example computing system upon which the various components, modules, systems, devices, and/or nodes may be implemented is provided. Various examples are provided throughout the specification.

(23) FIG. 2 is a logical block diagram illustrating an example of query processing groups performing queries to database data, according to some embodiments. Provider network **260** may be a private or closed system or may be set up by an entity such as a company or a public sector

organization to provide one or more services (such as various types of cloud-based storage) accessible via the Internet and/or other networks to clients **290**.

(24) Provider network **260** may be implemented in a single location or may include numerous data centers hosting various resource pools, such as collections of physical and/or virtualized computer servers, storage devices, networking equipment and the like (e.g., computing system **1000** described below with regard to FIG. **8**), needed to implement and distribute the infrastructure and storage services offered by the provider network **260**. The provider network **260** can be formed as a number of regions, where a region is a separate geographical area in which the cloud provider clusters data centers. Each region can include two or more availability zones connected to one another via a private high speed network, for example a fiber communication connection.

(25) An availability zone (also known as an availability domain, or simply a “zone”) refers to an isolated failure domain including one or more data center facilities with separate power, separate networking, and separate cooling from those in another availability zone. Preferably, availability zones within a region are positioned far enough away from one other that the same natural disaster should not take more than one availability zone offline at the same time. Customers can connect to availability zones of the provider network **260** via a publicly accessible network (e.g., the Internet, a cellular communication network). Regions are connected to a global network which includes private networking infrastructure (e.g., fiber connections controlled by the cloud provider) connecting each region to at least one other region. The provider network **260** may deliver content from points of presence outside of, but networked with, these regions by way of edge locations and regional edge cache servers. An edge location can be an extension of the cloud provider network outside of the traditional region/AZ context. For example an edge location can be a data center positioned to provide capacity to a set of customers within a certain latency requirement, a set of servers provided to a customer's premises, or a set of servers provided within (or forming part of) a cellular communications network, each of which can be controlled at least in part by the control plane of a nearby AZ or region. This compartmentalization and geographic distribution of computing hardware enables the provider network **260** to provide low-latency resource access to customers on a global scale with a high degree of fault tolerance and stability.

(26) The traffic and operations of the provider network may broadly be subdivided into two categories in various embodiments: control plane operations carried over a logical control plane and data plane operations carried over a logical data plane. While the data plane represents the movement of user data through the distributed computing system, the control plane represents the movement of control signals through the distributed computing system. The control plane generally includes one or more control plane components distributed across and implemented by one or more control servers. Control plane traffic generally includes administrative operations, such as system configuration and management (e.g., resource placement, hardware capacity management, diagnostic monitoring, system state information). The data plane includes customer resources that are implemented on the cloud provider network (e.g., computing instances, containers, block storage volumes, databases, file storage). Data plane traffic generally includes non-administrative operations such as transferring customer data to and from the customer resources. Certain control plane components (e.g., tier one control plane components such as the control plane for a virtualized computing service) are typically implemented on a separate set of servers from the data plane servers, while other control plane components (e.g., tier two control plane components such as analytics services) may share the virtualized servers with the data plane, and control plane traffic and data plane traffic may be sent over separate/distinct networks.

(27) In some embodiments, provider network **260** may implement various computing resources or services, such as database service(s) **280**, (e.g., relational database services, non-relational database services, a map reduce service, a data warehouse service, and/or other large scale data processing services or various other types database services), data storage service **270** (e.g., object storage services or block-based storage services that may implement a centralized data store for various

types of data), and/or any other type of network based services (which may include a virtual compute service and various other types of storage, processing, analysis, communication, event handling, visualization, and security services not illustrated).

(28) In various embodiments, the components illustrated in FIG. 2 may be implemented directly within computer hardware, as instructions directly or indirectly executable by computer hardware (e.g., a microprocessor or computer system), or using a combination of these techniques. For example, the components of FIG. 2 may be implemented by a system that includes a number of computing nodes (or simply, nodes), each of which may be similar to the computer system embodiment illustrated in FIG. 8 and described below. In various embodiments, the functionality of a given system or service component (e.g., a component of database service **280** or data storage service **270**) may be implemented by a particular node or may be distributed across several nodes. In some embodiments, a given node may implement the functionality of more than one service system component (e.g., more than one data store component).

(29) Database services **280** may be various types of data processing services that perform general or specialized data processing functions (e.g., anomaly detection, machine learning, data mining, big data querying, or any other type of data processing operation). For example, in at least some embodiments, database services **280** may include a map reduce service that creates clusters of processing nodes that implement map reduce functionality over data stored in the map reduce cluster as well as data stored in data storage service **270**. In another example, database service **280** may include various types of database services (both relational and non-relational) for storing, querying, and updating data. Such services may be enterprise-class database systems that are highly scalable and extensible. Queries may be directed to a database in data database service **280** that is distributed across multiple physical resources, and the database system may be scaled up or down on an as needed basis.

(30) Database service **280** may work effectively with database schemas of various types and/or organizations, in different embodiments. In some embodiments, clients/subscribers may submit queries in a number of ways, e.g., interactively via an SQL interface to the database system. In other embodiments, external applications and programs may submit queries using Open Database Connectivity (ODBC) and/or Java Database Connectivity (JDBC) driver interfaces to the database system. For instance, database service **280** may implement, in some embodiments, a data warehouse service, that utilizes another data processing service, to execute portions of queries or other access requests with respect to data that is stored in a remote data store, such as data storage service(s) **270** (or a data store external to provider network **260**) to implement distributed data processing for distributed data sets.

(31) In at least some embodiments, database service **280** may be a data warehouse service. Thus in the description that follows database service **280** may be discussed according to the various features or components that may be implemented as part of a data warehouse service, including control plane **250**, proxy service (not shown), and query processing groups **200**. Note that such features or components may also be implemented in a similar fashion for other types of database services and thus the following examples may be applicable to other types of database service **280**. Database service **280** may implement one (or more) processing clusters that are attached to a database (e.g., a data warehouse). In some embodiments, these processing clusters may be designated as a primary and secondary (or concurrent, additional, or burst processing clusters) that perform queries to an attached database warehouse.

(32) In embodiments where database service **280** is a data warehouse service, the data warehouse service may offer clients a variety of different data management services, according to their various needs. In some cases, clients may wish to store and maintain large of amounts data, such as sales records marketing, management reporting, business process management, budget forecasting, financial reporting, website analytics, or many other types or kinds of data. A client's use for the data may also affect the configuration of the data management system used to store the data. For

instance, for certain types of data analysis and other operations, such as those that aggregate large sets of data from small numbers of columns within each row, a columnar database table may provide more efficient performance. In other words, column information from database tables may be stored into data blocks on disk, rather than storing entire rows of columns in each data block (as in traditional database schemes). The following discussion describes various embodiments of a relational columnar database system implemented as a data warehouse. However, various versions of the components discussed below as may be equally adapted to implement embodiments for various other types of relational database systems, such as row-oriented database systems. Therefore, the following examples are not intended to be limiting as to various other types or formats of database systems.

(33) In some embodiments, storing table data in such a columnar fashion may reduce the overall disk I/O requirements for various queries and may improve analytic query performance. For example, storing database table information in a columnar fashion may reduce the number of disk I/O requests performed when retrieving data into memory to perform database operations as part of processing a query (e.g., when retrieving all of the column field values for all of the rows in a table) and may reduce the amount of data that needs to be loaded from disk when processing a query. Conversely, for a given number of disk requests, more column field values for rows may be retrieved than is necessary when processing a query if each data block stored entire table rows. In some embodiments, the disk requirements may be further reduced using compression methods that are matched to the columnar storage data type. For example, since each block contains uniform data (i.e., column field values that are all of the same data type), disk storage and retrieval requirements may be further reduced by applying a compression method that is best suited to the particular column data type. In some embodiments, the savings in space for storing data blocks containing only field values of a single column on disk may translate into savings in space when retrieving and then storing that data in system memory (e.g., when analyzing or otherwise processing the retrieved data).

(34) Database service **280** may be implemented by a large collection of computing devices, such as customized or off-the-shelf computing systems, servers, or any other combination of computing systems or devices, such as the various types of systems **1000** described below with regard to FIG. **8**. Different subsets of these computing devices may be controlled by control plane **250**. Control plane **250**, for example, may provide a control interface to clients or users who wish to interact with the processing groups, such as query processing groups **200** managed by control plane **250**. For example, control plane **250** may generate one or more graphical user interfaces (GUIs) for clients, which may then be utilized to select various control functions offered by the control interface for the query processing groups **200** hosted in the database service **280**. Control plane **250** may provide or implement access to various metrics collected for the performance of different features of database service **280**, including processing cluster performance, in some embodiments.

(35) As discussed above, various clients (or customers, organizations, entities, or users) may wish to store and manage data using a database service **280**. Query processing groups **200** may respond to various requests, including write/update/store requests (e.g., to write data into storage) or queries for data (e.g., such as a Server Query Language request (SQL) for particular data). For example, multiple users or clients may access a processing cluster to obtain data warehouse services.

(36) For databases manually managed by users, database service **280** may provide network endpoints directly to the clusters which allow the users manage in order to implement client applications that send requests and other messages directly to a particular cluster. Network endpoints, for example may be a particular network address, such as a URL, which points to a particular cluster. For instance, a client may be given the network endpoint “http://mycluster.com” to send various request messages to. Multiple clients (or users of a particular client) may be given a network endpoint for a particular cluster. Various security features may be implemented to prevent unauthorized users from accessing the clusters.

(37) In at least some embodiments, database service **280** may implement proxy service to provide access to databases (e.g., data warehouses) hosted in database service **280**. For databases managed by database service **280**, database service **280** may provide database endpoints (e.g., network endpoints) for a hosted database. Database endpoints may not provide direct access to a particular query processing group **200**, as the processing group used to respond to such requests (e.g., queries) may change according to various scaling techniques. Instead, client applications may utilize the database endpoint for a database to be included in various client applications or other communications for database access so that proxy service can direct the requests to the appropriate processing cluster without the client application having to be altered every time a change in processing cluster (e.g., scaling operations) are performed by database service **280**. In this way, database service **280** can perform scaling and other management operations without interfering with client applications.

(38) Query processing groups **200** hosted by database service **280** may provide an enterprise-class database query and management system that allows users to send data processing requests to be executed by the groups **200**, such as by sending a query. Query processing groups **200** may perform data processing operations with respect to data stored locally in a processing cluster, as well as remotely stored data. For example, data storage service **270** implemented by provider network **260** that stores remote data, such as backups or other data of a database stored in a cluster. In some embodiments, database data **240** may not be stored locally in a query processing group **200** but instead may be stored in data storage service **270** (e.g., with data being partially or temporarily stored in query processing group **200** to perform queries). Queries sent to a processing cluster **23** (or routed/redirect/assigned/allocated to query processing group(s)) may be directed to local data stored in the processing cluster and/or remote data. Therefore, processing clusters may implement local data processing, such as local data processing, to plan and execute the performance of queries with respect to local data in the processing cluster, as well as a remote data processing client.

(39) Database service **280** may implement different types or configurations of processing clusters. For example, different configurations may utilize various different configurations of computing resources, including, but not limited to, different numbers of computational nodes, different processing capabilities (e.g., processor size, power, custom or task-specific hardware, such as hardware accelerators to perform different operations, such as regular expression searching or other data processing operations), different amounts of memory, different networking capabilities, and so on. Thus, for some queries, different configurations of query processing groups **200** may offer different execution times. Different configurations of query processing groups **200** may be maintained in different pools of available processing clusters to be attached to a database. Attached processing clusters may then be made exclusively assigned or allocated for the use of performing queries to the attached database, in some embodiments. The number of query processing groups **200** attached to a database may change over time according to the selection techniques discussed below.

(40) In some embodiments, database service **280** may have at least one query processing group attached to a database, which may be the “primary group.” Primary groups may be reserved, allocated, permanent, or otherwise dedicated processing resources that store and/or provide access to a database for a client, in some embodiments. Primary groups, however, may be changed. For example, a different processing cluster may be attached to a database and then designated as the primary database (e.g., allowing an old primary group to still be used as a “secondary” processing cluster or released to a pool of processing clusters made available to be attached to a different database). Techniques to resize or change to a different configuration of a primary group may be performed, in some embodiments. The available processing clusters that may also be attached, as determined, to a database may be maintained (as noted earlier) in different configuration type pools, which may be a set of warmed, pre-configured, initialized, or otherwise prepared clusters which may be on standby to provide additional query performance capacity in addition to that

provided by a primary group. Control plane **250** may manage cluster pools by managing the size of cluster pools (e.g., by adding or removing processing clusters based on demand to use the different processing clusters).

(41) As databases are created, updated, and/or otherwise modified, snapshots, copies, or other replicas of the database at different states may be stored separate from database service **280** in data storage service **250**, in some embodiments. For example, a leader node, or other processing cluster component, may implement a backup agent or system that creates and store database backups for a database to be stored as database data **240** in data storage service **270**. Database data **240** may include user data (e.g., tables, rows, column values, etc.) and database metadata (e.g., information describing the tables which may be used to perform queries to a database, such as schema information, data distribution, range values or other content descriptors for filtering out portions of a table from a query, a superblock, etc.). A timestamp or other sequence value indicating the version of database data **240** may be maintained in some embodiments, so that the latest database data **240** may, for instance, be obtained by a processing cluster in order to perform queries. In at least some embodiments, database data **240** may be treated as the authoritative version of data, and data stored in query processing groups **200** for local processing as a cached version of data.

(42) Data storage service **270** may implement different types of data stores for storing, accessing, and managing data on behalf of clients **250** as a network-based service that enables clients **250** to operate a data storage system in a cloud or network computing environment. Data storage service(s) **270** may also include various kinds of object or file data stores for putting, updating, and getting data objects or files. For example, one data storage service **270** may be an object-based data store that allows for different data objects of different formats or types of data, such as structured data (e.g., database data stored in different database schemas), unstructured data (e.g., different types of documents or media content), or semi-structured data (e.g., different log files, human-readable data in different formats like JavaScript Object Notation (JSON) or Extensible Markup Language (XML)) to be stored and managed according to a key value or other unique identifier that identifies the object. In at least some embodiments, data storage service(s) **270** may be treated as a data lake. For example, an organization may generate many different kinds of data, stored in one or multiple collections of data objects in a data storage service **270**. The data objects in the collection may include related or homogenous data objects, such as database partitions of sales data, as well as unrelated or heterogeneous data objects, such as audio files and web site log files. Data storage service(s) **270** may be accessed via programmatic interfaces (e.g., APIs) or graphical user interfaces. For example, a format independent data processing service may access data objects stored in data storage services via the programmatic interfaces.

(43) Generally speaking, clients **250** may encompass any type of client that can submit network-based requests to provider network **260** via a network, including requests for storage services (e.g., a request to query a database service **280**, or a request to create, read, write, obtain, or modify data in data storage service(s) **270**, etc.). For example, a given client **250** may include a suitable version of a web browser, or may include a plug-in module or other type of code module that can execute as an extension to or within an execution environment provided by a web browser. Alternatively, a client **250** may encompass an application such as a database application (or user interface thereof), a media application, an office application or any other application that may make use of database service(s) **280** or storage resources in data storage service(s) **270** to store and/or access the data to implement various applications. In some embodiments, such an application may include sufficient protocol support (e.g., for a suitable version of Hypertext Transfer Protocol (HTTP)) for generating and processing network-based services requests without necessarily implementing full browser support for all types of network-based data. That is, client **250** may be an application that can interact directly with provider network **260**. In some embodiments, client **250** may generate network-based services requests according to a Representational State Transfer (REST)-style network-based services architecture, a document- or message-based network-based services

architecture, or another suitable network-based services architecture.

(44) In some embodiments, a client **250** may provide access to provider network **260** to other applications in a manner that is transparent to those applications. For example, client **250** may integrate with an operating system or file system to provide storage on one of data storage service(s) **270** (e.g., a block-based storage service). However, the operating system or file system may present a different storage interface to applications, such as a conventional file system hierarchy of files, directories and/or folders. In such an embodiment, applications may not need to be modified to make use of the storage system service model. Instead, the details of interfacing to the data storage service(s) **270** may be coordinated by client **250** and the operating system or file system on behalf of applications executing within the operating system environment. Similarly, a client **250** may be an analytics application that relies upon data processing service(s) to execute various queries for data already ingested or stored in the data processing service (e.g., such as data maintained in a data warehouse service).

(45) Clients **250** may convey network-based services requests (e.g., access requests to read or write data may be directed to data in data storage service(s) **270**, or operations, tasks, or jobs, such as queries, being performed as part of data processing service(s)) to and receive responses from provider network **260** via network **260**. In various embodiments, network **260** may encompass any suitable combination of networking hardware and protocols necessary to establish network-based communications between clients **250** and provider network **260**. For example, network **260** may generally encompass the various telecommunications networks and service providers that collectively implement the Internet. Network **260** may also include private networks such as local area networks (LANs) or wide area networks (WANs) as well as public or private wireless networks. For example, both a given client **250** and provider network **260** may be respectively provisioned within enterprises having their own internal networks. In such an embodiment, network **260** may include the hardware (e.g., modems, routers, switches, load balancers, proxy servers, etc.) and software (e.g., protocol stacks, accounting software, firewall/security software, etc.) necessary to establish a networking link between given client **250** and the Internet as well as between the Internet and provider network **260**. It is noted that in some embodiments, clients **250** may communicate with provider network **260** using a private network rather than the public Internet. In some embodiments, clients of data processing services and/or data storage service(s) **270** may be implemented within provider network **260** (e.g., an application hosted on a virtual computing resource that utilizes a data processing service to perform database queries) to implement various application features or functions and thus various features of client(s) **250** discussed above may be applicable to such internal clients as well.

(46) As illustrated in the example of FIG. 2, a query processing group **200** may include a leader node **210** and compute node(s) **220**, which may communicate with each other over an interconnect (not illustrated). Leader node **210** may implement query planning **212** to generate query plan(s), query execution **214** for executing queries on query processing group **200** that perform data processing that can utilize remote query processing resources for remotely stored data (e.g., by utilizing one or more query execution slot(s)/queue(s) **217**). As described herein, each node in a processing cluster **200** may include attached storage, such as attached storage **222** on which a database (or portions thereof) may be stored on behalf of clients (e.g., users, client applications, and/or storage service subscribers).

(47) Note that in at least some embodiments, query processing capability may be separated from compute nodes, and thus in some embodiments, additional components may be implemented for processing queries. Additionally, it may be that in some embodiments, no one node in processing cluster **200** is a leader node as illustrated in FIG. 2, but rather different nodes of the nodes in processing cluster **200** may act as a leader node or otherwise direct processing of queries to data stored in processing cluster **200**. While nodes of processing cluster may be implemented on separate systems or devices, in at least some embodiments, some or all of processing cluster may

be implemented as separate virtual nodes or instance on the same underlying hardware system (e.g., on a same server).

(48) Leader node **210** may manage communications with clients **290**. Leader node **710** may communicate with proxy service (not shown) and may receive query **201** and return query results **203** to proxy service (instead of communicating directly with a client application).

(49) Leader node **210** may be a node that receives a query **201** from various client programs (e.g., applications) and/or subscribers (users) (either directly or routed to leader node **210** from proxy service), then parses them and develops an execution plan (e.g., query plan(s)) to carry out the associated database operation(s)). More specifically, leader node **210** may develop the series of steps necessary to obtain results for the query. Query **201** may be directed to data that is stored both locally within processing cluster **200** (e.g., at one or more of compute nodes **220**) and data stored remotely. Leader node **210** may also manage the communications among compute nodes **220** instructed to carry out database operations for data stored in the processing cluster **200**. For example, node-specific query instructions **204** may be generated or compiled code by query execution **214** that is distributed by leader node **210** to various ones of the compute nodes **220** to carry out the steps needed to perform query **201**, including executing the code to generate intermediate results of query **201** at individual compute nodes may be sent back to the leader node **210**. Leader node **210** may receive data and query responses or results from compute nodes **220** in order to determine a final result **203** for query **201**.

(50) A database schema, data format and/or other metadata information for the data stored among the compute nodes, such as the data tables stored in the cluster, may be managed and stored by leader node **210**. Query planning **212** may account for remotely stored data by generating node-specific query instructions that include remote operations to be directed by individual compute node(s). Although not illustrated, in some embodiments, a leader node may implement burst manager to send **206** a query plan generated by query planning **212** to be performed at another attached processing cluster and return results **208** received from the burst processing cluster to a client as part of results **203**.

(51) In at least some embodiments, a result cache **219** may be implemented as part of leader node **210**. For example, as query results are generated, the results may also be stored in result cache **219** (or pointers to storage locations that store the results either in primary processing cluster **200** or in external storage locations), in some embodiments. Result cache **219** may be used instead of other processing cluster capacity, in some embodiments, by recognizing queries which would otherwise be sent to another attached processing cluster to be performed that have results stored in result cache **219**. Various caching strategies (e.g., LRU, FIFO, etc.) for result cache **219** may be implemented, in some embodiments. Although not illustrated in FIG. 2, result cache **219** could be stored in other storage systems (e.g., other storage services, such as a NoSQL database) and/or could store sub-query results.

(52) Processing cluster **200** may also include compute nodes **220**. Compute nodes **220**, may for example, be implemented on servers or other computing devices, such as those described below with regard to computer system **1000** in FIG. 8, and each may include individual query processing “slices” defined, for example, for each core of a server's multi-core processor, one or more query processing engine(s), such as query engine(s) **224**, to execute the instructions **204** or otherwise perform the portions of the query plan assigned to the compute node. Query engine(s) **224** may access a certain memory and disk space in order to process a portion of the workload for a query (or other database operation) that is sent to one or more of the compute nodes **220**. Query engine **224** may access attached storage, such as **222**, to perform local operation(s), such as local operations **218**. For example, query engine **224** may scan data in attached storage **222**, access indexes, perform joins, semi joins, aggregations, or any other processing operation assigned to the compute node **220**.

(53) Query engine **224** may also direct the execution of remote data processing operations, by

providing remote operation(s), such as remote operations **216**, to remote data processing clients, such as remote data processing client **226**. Remote data processing clients **226** may be implemented by a client library, plugin, driver or other component that sends request sub-queries to be performed by data storage service **270** or requests to for data, **232**. As noted above, in some embodiments, Remote data processing clients **226** may read, process, or otherwise obtain data **234** in response from database data **240** in data storage service **270**, which may further process, combine, and or include them with results of location operations **218**.

(54) Compute nodes **220** may send intermediate results from queries back to leader node **210** for final result generation (e.g., combining, aggregating, modifying, joining, etc.). Remote data processing clients **226** may retry data requests **232** that do not return within a retry threshold.

(55) Attached storage **222** may be implemented as one or more of any type of storage devices and/or storage system suitable for storing data accessible to the compute nodes, including, but not limited to: redundant array of inexpensive disks (RAID) devices, disk drives (e.g., hard disk drives or solid state drives) or arrays of disk drives such as Just a Bunch Of Disks (JBOD), (used to refer to disks that are not implemented according to RAID), optical storage devices, tape drives, RAM disks, Storage Area Network (SAN), Network Access Storage (NAS), or combinations thereof. In various embodiments, disks may be formatted to store database tables (e.g., in column oriented data formats or other data formats).

(56) FIG. **3A** is a logical block diagram of a conventional shared memory processing system, according to some embodiments. Processing groups **320**, including groups **320a-320m**, may collectively share a global shared memory **300**. Data objects may be allocated from individual shared memory pages **310**, including pages **310a**, **310b** and **310c-310n** to produce allocated objects **321**, including allocated objects **321a-321m**, in various embodiments. Memory allocation and deallocation requests **330** may be performed between the various processing groups **320** and the global shared memory **300**.

(57) In a conventional shared memory processing system, individual objects may be allocated and deallocated in the global shared memory pool based directly and may distribute the objects to different physical pages. This prevents the identifying of affected memory pages upon abnormal execution of a process of a processing group. Furthermore, this precludes the use of processor protection flags to protect the objects from unintended changes. Thus, abnormal execution of a process may lead to corruption of global memory state in a conventional shared memory processing system.

(58) FIG. **3B** is a logical block diagram of a shared memory processing system with improved crash resilience, according to some embodiments. Processing groups **320**, such as the processing groups **120** of FIG. **1**, may collectively share a global shared memory **300**, such as provided by the shared memory manager **130** of FIG. **1**, in various embodiments. Rather than allocate objects from pages of the global shared memory directly, memory allocation and deallocation requests **330** may instead be performed between the various processing groups **320** and the memory arenas **340** and **341**, including private arenas **340a-340m** and public arena **341**. Examples of such arenas are the public arena **134** and private arena(s) **136** as shown above in FIG. **1**. Pages **310** of the global shared memory **300** may be allocated as needed to the various memory arenas.

(59) Individual ones of the processing groups **320** may have an associated group-private arena **340** for which the processing group has read-write access. Modifications to data objects allocated within a group-private arena **340** may only affect processing state of the individual processing group. Therefore, abnormal execution of a process of a processing group may only affect that processing group provided that modified objects are constrained only to the group-private arena **340**, in some embodiments. In these cases, recovery of abnormal execution of a process of a processing group may be performed without affecting other processing groups, in some embodiments. A processing group is excluded from access to group-private arenas **340** that are associated with other processing groups different from the processing group.

(60) In order to support allocating memory for truly shared objects such as disk blocks. IPC messages, communications packets, logging and the like, a public arena **341** may also be provided. Unlike the group-private arenas **340**, the public arena **341** may be accessible by all processing groups **320**. However, individual processing groups **320** may, by default, have read-only access to data objects allocated in the public arena **341**. Thus, abnormal execution of a process of a processing group may still avoid corruption of shared memory state, in some embodiments. In order to affect changes to data objects in the public arena, a process of a processing group **320** may elevate execution privilege in order to obtain read-write access to object(s) in the public arena **341**, as discussed further below in FIG. 7. During this time of elevated privilege, abnormal execution of the process may lead to corruption of global memory state, in some embodiments.

(61) By organizing allocation of data objects into various arenas, control over allocation of data objection into specific memory pages **310** may be accomplished, in various embodiments. This enables the use of per page memory protection featured of modern computing processors and memory management units, this leading to more efficient crash resiliency in shared memory data processing systems, in some embodiments.

(62) FIG. 4 is a flow diagram illustrating improved crash resilience, according to some embodiments. In data processing application implementing improved crash resiliency, as shown in **400**, a processing group such as the processing group **120** of FIG. 1, including multiple executing processes may allocate shared memory from a global shared memory pool, such as the memory pool **132** of FIG. 1, into a group-private shared memory arena, such as the private arena **136** of FIG. 1, in some embodiments. As discussed above in FIG. 3, processes of the processing group may have read-write access to various data objects allocated within this allocated memory in the group-private shared memory arena, in some embodiments.

(63) As shown in **410**, one or more of the processes of the processing group may employ the read-write access to the allocated memory to perform one or more updates to data objects allocated within the allocated memory in the group-private shared memory arena.

(64) As shown in **420**, an abnormal execution condition, such as a protected memory fault, an uncaught processing exception or an assertion from third party software, may then be detected at a failing process of the processing group, in some embodiments. This abnormal execution condition may be detected during execution of a processing job assigned to the processing group, in various embodiments.

(65) Responsive to this detection, the failing process, as discussed further in FIGS. 7 and 8 below, may determine if it is a leader process, as shown in **430**. If the failing process is a leader process, such as the leader process **122** of FIG. 1, as indicated in a positive exit at **430**, the method may proceed to step **440**. If the failing process is not a leader process but a compute process, such as the compute process **124** of FIG. 1, as indicated in a negative exit at **430**, the method may proceed to step **435**.

(66) As shown in step **435**, if the failing process is not a leader process, the failing process may then send a message to the leader process of the processing group indicating the detected abnormal execution condition, in some embodiments. Upon receiving this message, the leader process may then gracefully terminate the failing process, as shown in **436**. The method may then proceed to step **440**.

(67) As shown in step **440**, if the failing process is a leader process, the leader process may then abort the assigned processing job and restart all processes, as shown in **450**, of the processing group while continuing execution of processes in other processing groups, in some embodiments. By restarting only the processing group without restarting other processes in other groups, disruption of the data processing application may be minimized, in some embodiments. The method may then proceed to step **460**.

(68) As shown in **460**, the method may then return the allocated shared memory in the group-private memory arena to the global shared memory, in some embodiments.

(69) FIGS. 5A and 5B are logical block diagrams illustrating fault signaling in a shared memory processing system with improved crash resilience, according to some embodiments. One or more computing nodes, may implement a compute node **500** and a leader node **540**, in some embodiments. In some embodiments, these nodes may be implemented on a same computing system, such as described below in FIG. 8, while in other embodiments, the nodes may be implemented using separate systems to form a distributed system.

(70) As shown in FIG. 5A, a computing node **500** may implement at least one compute process **510**, a process dispatcher **520** and a resource manager **530**, in some embodiments. The compute process **510** may implement one or more execution threads (not shown) with each execution thread implementing a signal handler **512**. These signal handlers may respectively communicate via a pipe **514** to a cleanup thread **516**, in some embodiments.

(71) If an abnormal execution condition, such as a protected memory fault, an uncaught processing exception or an assertion from third party software, is detected at a signal handler **512**, the signal handler may write information describing the abnormal execution condition to the pipe **514**. This may allow the signal handler **514** to communicate information regarding the uncaught processing exception while complying with the execution constraints of an exception or signal handler. Once the information is written, the signal handler may enter an infinite execution loop until the compute process **510** terminates, in some embodiments.

(72) Upon receiving abnormal execution condition information from the pipe **514**, the cleanup thread **516** may send a taskAborted message **504** to the leader node **540** of the process group. The leader node **540** may then send cmdAbort message(s) **502** to various ones of the computing nodes of the process group, in some embodiments.

(73) Should the dispatcher **520** receive a cmdAbort message **502** from the leader process **540**, the dispatcher may determine if a shared memory state may be safely recovered. In some embodiments, to make this determination, the dispatcher **520** may determine if the compute process **510** possesses elevated privileges to modify public shared memory, in some embodiments. This elevated privilege may be indicated in a number of ways. For example, compute process privilege may be indicated by explicit memory access privileges in some embodiments while in other embodiments compute process privilege may be indicated through control of synchronization primitives such as lock or mutex ownership. These examples are not intended to be limiting, however, and any number of compute process privilege indications may be envisioned.

(74) Should elevated compute process privilege be determined, the dispatcher **520** may indicate a need to restart the entire application including all processes of all processing groups, in some embodiments. Should no compute process privilege be determined, the dispatcher **520** may send a refork request for the compute process **510** to the resource manager **530** after removing the compute process from the processing group. The dispatcher **520** may then send a SIGKILL signal to terminated to the compute process **510**, in some embodiments.

(75) As shown in FIG. 5B, the leader process may implement one or more execution threads (not shown) with each execution thread implementing a signal handler **552**. These signal handlers may respectively communicate via a pipe **554** to a cleanup thread **556**, in some embodiments.

(76) If an abnormal execution condition, such as a protected memory fault, an uncaught processing exception or an assertion from third party software, is detected at a signal handler **552**, the signal handler may write information describing the abnormal execution condition to the pipe **554**. This may allow the signal handler **554** to communicate information regarding the uncaught processing exception while complying with the execution constraints of an exception or signal handler. Once the information is written, the signal handler may enter an infinite execution loop until the leader process **550** terminates, in some embodiments.

(77) Upon receiving abnormal execution condition information from the pipe **554**, the cleanup thread **556** may send cmdAbort messages **502** to various ones of the computing nodes of the process group, in some embodiments. As a result, individual ones of the computing nodes will

proceed as discussed above with respect to FIG. 5A, in some embodiments. The cleanup thread **556** may then send a refork request to a resource manager **560** executing on the leader node after removing the compute process from the processing group. The resource manager **560** may then send a SIGKILL signal to terminated to the compute process **560**, in some embodiments.

(78) FIG. **6** is a flow diagram illustrating fault signaling in a shared memory processing system with improved crash resilience, according to some embodiments. In data processing application implementing improved crash resiliency, as shown in **600**, a processing group such as the processing group **120** of FIG. **1**, including multiple executing processes may allocate shared memory from a global shared memory pool, such as the memory pool **132** of FIG. **1**, into a group-private shared memory arena, such as the private arena **136** of FIG. **1**, in some embodiments. As discussed above in FIG. **3**, processes of the processing group may have read-write access to various data objects allocated within this allocated memory in the group-private shared memory arena, in some embodiments.

(79) As shown in **610**, one or more of the processes of the processing group may employ the read-write access to the allocated memory to perform one or more updates to data objects allocated within the allocated memory in the group-private shared memory arena.

(80) As shown in **620**, an abnormal execution condition, such as a protected memory fault, an uncaught processing exception or an assertion from third party software, may then be detected at a failing process of the processing group, in some embodiments. This abnormal execution condition may be detected during execution of a processing job assigned to the processing group, in various embodiments.

(81) Responsive to this detection, the failing process may determine if it is a leader process, as shown in **630**. If the failing process is a leader process, such as the leader process **122** of FIG. **1**, as indicated in a positive exit at **630**, the method may proceed to step **640**. If the failing process is not a leader process but a compute process, such as the compute process **124** of FIG. **1**, as indicated in a negative exit at **630**, the method may proceed to step **635**.

(82) As shown in step **635**, if the failing process is not a leader process, the failing process may then send a message to the leader process of the processing group indicating the detected abnormal execution condition, in some embodiments. Upon receiving this message, the leader process may then gracefully terminate the failing process, as shown in **436**. The method may then proceed to step **440**.

(83) As shown in step **640**, if the failing process is a leader process, the leader process may then abort the assigned processing job and restart all processes, as shown in **650**, of the processing group while continuing execution of processes in other processing groups, in some embodiments. By restarting only the processing group without restarting other processes in other groups, disruption of the data processing application may be minimized, in some embodiments. The method may then proceed to step **660**.

(84) As shown in **660**, the method may then return the allocated shared memory in the group-private memory arena to the global shared memory, in some embodiments.

(85) FIG. **7** is a flow diagram illustrating updating global memory state in a shared memory processing system with improved crash resilience, according to some embodiments. The method begins at step **700** where memory for a shared object may be allocated in a public shared memory arena from a global shared memory, in some embodiments. Examples of such shared objects are objects that are shared across multiple queries including network packets (packets headers), database row sets, log files, access credentials, results caches, code pools and so forth. Such examples are not intended to be limiting and any number of shared objects may be envisioned.

(86) Once a determination has been made to modify the public shared memory object, as shown in **710**, access privilege to the public shared memory object may be elevated, in some embodiments, to allow writing privilege to the allocated shared memory. For example, privilege may be elevated through an explicit memory access privilege escalation in some embodiments while in other

embodiments privilege may be elevated through allocation of synchronization primitives such as a lock or mutex. These examples are not intended to be limiting, however, and any number of privilege elevation mechanisms may be envisioned.

(87) Once access privilege has been elevated, in some embodiments the allocated shared memory may be modified and committed, as shown in **720**. During this process, it should be noted that any abnormal execution condition of a process with such elevated privilege may result in a corrupted state of global shared memory.

(88) Upon completion of modification of the allocated memory, as shown in **730**, access privilege to the allocated shared memory may then be lowered, in some embodiments. Upon lowering of access privilege, an abnormal execution condition of a process with such lowered privilege may occur without corruption of a state of global shared memory.

(89) The methods described herein may in various embodiments be implemented by any combination of hardware and software. For example, in one embodiment, the methods may be implemented by a computer system (e.g., a computer system as in FIG. **8**) that includes one or more processors executing program instructions stored on a computer-readable storage medium coupled to the processors. The program instructions may implement the functionality described herein (e.g., the functionality of various servers and other components that implement the network-based virtual computing resource provider described herein). The various methods as illustrated in the figures and described herein represent example embodiments of methods. The order of any method may be changed, and various elements may be added, reordered, combined, omitted, modified, etc.

(90) Various ones of the illustrated embodiments may include one or more computer systems **1000** such as that illustrated in FIG. **8** or one or more components of the computer system **1000** that function in a same or similar way as described for the computer system **1000**.

(91) In the illustrated embodiment, computer system **1000** includes one or more processors **1010** coupled to a system memory **1020** via an input/output (I/O) interface **1030**. Computer system **1000** further includes a network interface **1040** coupled to I/O interface **1030**. In some embodiments, computer system **1000** may be illustrative of servers implementing enterprise logic or downloadable applications, while in other embodiments servers may include more, fewer, or different elements than computer system **1000**.

(92) Computer system **1000** includes one or more processors **1010** (any of which may include multiple cores, which may be single or multi-threaded) coupled to a system memory **1020** via an input/output (I/O) interface **1030**. Computer system **1000** further includes a network interface **1040** coupled to I/O interface **1030**. In various embodiments, computer system **1000** may be a uniprocessor system including one processor **1010**, or a multiprocessor system including several processors **1010** (e.g., two, four, eight, or another suitable number). Processors **1010** may be any suitable processors capable of executing instructions. For example, in various embodiments, processors **1010** may be general-purpose or embedded processors implementing any of a variety of instruction set architectures (ISAs), such as the x86, PowerPC, SPARC, or MIPS ISAs, or any other suitable ISA. In multiprocessor systems, each of processors **1010** may commonly, but not necessarily, implement the same ISA. The computer system **1000** also includes one or more network communication devices (e.g., network interface **1040**) for communicating with other systems and/or components over a communications network (e.g. Internet, LAN, etc.). For example, a client application executing on system **1000** may use network interface **1040** to communicate with a server application executing on a single server or on a cluster of servers that implement one or more of the components of the embodiments described herein. In another example, an instance of a server application executing on computer system **1000** may use network interface **1040** to communicate with other instances of the server application (or another server application) that may be implemented on other computer systems.

(93) System memory **1010** may store instructions and data accessible by processor **1010**. In various

embodiments, system memory **1020** may be implemented using any suitable memory technology, such as static random-access memory (SRAM), synchronous dynamic RAM (SDRAM), non-volatile/Flash-type memory, or any other type of memory. In the illustrated embodiment, program instructions and data implementing desired functions, such as those methods and techniques as described above for providing a database compute node, as indicated at **1026**, are shown stored within system memory **1020** as program instructions **1025**. In some embodiments, system memory **1020** may include data store **1045** which may be configured as described herein.

(94) In some embodiments, system memory **1020** may be one embodiment of a computer-accessible medium that stores program instructions and data as described above. However, in other embodiments, program instructions and/or data may be received, sent or stored upon different types of computer-accessible media. Generally speaking, a computer-accessible medium may include computer-readable storage media or memory media such as magnetic or optical media, e.g., disk or DVD/CD-ROM coupled to computer system **1000** via I/O interface **1030**. A computer-readable storage medium may also include any volatile or non-volatile media such as RAM (e.g. SDRAM, DDR SDRAM, RDRAM, SRAM, etc.), ROM, etc., that may be included in some embodiments of computer system **1000** as system memory **1020** or another type of memory. Further, a computer-accessible medium may include transmission media or signals such as electrical, electromagnetic, or digital signals, conveyed via a communication medium such as a network and/or a wireless link, such as may be implemented via network interface **1040**.

(95) In one embodiment, I/O interface **1030** may coordinate I/O traffic between processor **1010**, system memory **1020** and any peripheral devices in the system, including through network interface **1040** or other peripheral interfaces. In some embodiments, I/O interface **1030** may perform any necessary protocol, timing or other data transformations to convert data signals from one component (e.g., system memory **1020**) into a format suitable for use by another component (e.g., processor **1010**). In some embodiments, I/O interface **1030** may include support for devices attached through various types of peripheral buses, such as a variant of the Peripheral Component Interconnect (PCI) bus standard or the Universal Serial Bus (USB) standard, for example. In some embodiments, the function of I/O interface **1030** may be split into two or more separate components, such as a north bridge and a south bridge, for example. Also, in some embodiments, some or all of the functionality of I/O interface **1030**, such as an interface to system memory **1020**, may be incorporated directly into processor **1010**.

(96) Network interface **1040** may allow data to be exchanged between computer system **1000** and other devices attached to a network, such as between a client device and other computer systems, or among hosts, for example. In particular, network interface **1040** may allow communication between computer system **800** and/or various other device **1060** (e.g., I/O devices). Other devices **1060** may include scanning devices, display devices, input devices and/or other communication devices, as described herein. Network interface **1040** may commonly support one or more wireless networking protocols (e.g., Wi-Fi/IEEE 802.7, or another wireless networking standard). However, in various embodiments, network interface **1040** may support communication via any suitable wired or wireless general data networks, such as other types of Ethernet networks, for example. Additionally, network interface **1040** may support communication via telecommunications/telephony networks such as analog voice networks or digital fiber communications networks, via storage area networks such as Fibre Channel SANs, or via any other suitable type of network and/or protocol.

(97) In some embodiments, I/O devices may be relatively simple or “thin” client devices. For example, I/O devices may be implemented as dumb terminals with display, data entry and communications capabilities, but otherwise little computational functionality. However, in some embodiments, I/O devices may be computer systems implemented similarly to computer system **1000**, including one or more processors **1010** and various other devices (though in some embodiments, a computer system **1000** implementing an I/O device **1050** may have somewhat

different devices, or different classes of devices).

(98) In various embodiments, I/O devices (e.g., scanners or display devices and other communication devices) may include, but are not limited to, one or more of: handheld devices, devices worn by or attached to a person, and devices integrated into or mounted on any mobile or fixed equipment, according to various embodiments. I/O devices may further include, but are not limited to, one or more of: personal computer systems, desktop computers, rack-mounted computers, laptop or notebook computers, workstations, network computers, “dumb” terminals (i.e., computer terminals with little or no integrated processing ability), Personal Digital Assistants (PDAs), mobile phones, or other handheld devices, proprietary devices, printers, or any other devices suitable to communicate with the computer system **1000**. In general, an I/O device (e.g., cursor control device, keyboard, or display(s)) may be any device that can communicate with elements of computing system **1000**.

(99) The various methods as illustrated in the figures and described herein represent illustrative embodiments of methods. The methods may be implemented manually, in software, in hardware, or in a combination thereof. The order of any method may be changed, and various elements may be added, reordered, combined, omitted, modified, etc. For example, in one embodiment, the methods may be implemented by a computer system that includes a processor executing program instructions stored on a computer-readable storage medium coupled to the processor. The program instructions may be configured to implement the functionality described herein.

(100) Various modifications and changes may be made as would be obvious to a person skilled in the art having the benefit of this disclosure. It is intended to embrace all such modifications and changes and, accordingly, the above description to be regarded in an illustrative rather than a restrictive sense.

(101) Various embodiments may further include receiving, sending or storing instructions and/or data implemented in accordance with the foregoing description upon a computer-accessible medium. Generally speaking, a computer-accessible medium may include storage media or memory media such as magnetic or optical media, e.g., disk or DVD/CD-ROM, volatile or non-volatile media such as RAM (e.g. SDRAM, DDR, RDRAM, SRAM, etc.), ROM, etc., as well as transmission media or signals such as electrical, electromagnetic, or digital signals, conveyed via a communication medium such as network and/or a wireless link.

(102) Embodiments of decentralized application development and deployment as described herein may be executed on one or more computer systems, which may interact with various other devices. FIG. 6 is a block diagram illustrating an example computer system, according to various embodiments. For example, computer system **1000** may be configured to implement nodes of a compute cluster, a distributed key value data store, and/or a client, in different embodiments. Computer system **1000** may be any of various types of devices, including, but not limited to, a personal computer system, desktop computer, laptop or notebook computer, mainframe computer system, handheld computer, workstation, network computer, a consumer device, application server, storage device, telephone, mobile telephone, or in general any type of compute node, computing node, or computing device.

(103) In the illustrated embodiment, computer system **1000** also includes one or more persistent storage devices **1060** and/or one or more I/O devices **1080**. In various embodiments, persistent storage devices **1060** may correspond to disk drives, tape drives, solid state memory, other mass storage devices, or any other persistent storage device. Computer system **1000** (or a distributed application or operating system operating thereon) may store instructions and/or data in persistent storage devices **1060**, as desired, and may retrieve the stored instruction and/or data as needed. For example, in some embodiments, computer system **1000** may be a storage host, and persistent storage **1060** may include the SSDs attached to that server node.

(104) In some embodiments, program instructions **1025** may include instructions executable to implement an operating system (not shown), which may be any of various operating systems, such

as UNIX, LINUX, Solaris™, MacOS™, Windows™, etc. Any or all of program instructions **1025** may be provided as a computer program product, or software, that may include a non-transitory computer-readable storage medium having stored thereon instructions, which may be used to program a computer system (or other electronic devices) to perform a process according to various embodiments. A non-transitory computer-readable storage medium may include any mechanism for storing information in a form (e.g., software, processing application) readable by a machine (e.g., a computer). Generally speaking, a non-transitory computer-accessible medium may include computer-readable storage media or memory media such as magnetic or optical media, e.g., disk or DVD/CD-ROM coupled to computer system **1000** via I/O interface **1030**. A non-transitory computer-readable storage medium may also include any volatile or non-volatile media such as RAM (e.g. SDRAM, DDR SDRAM, RDRAM, SRAM, etc.), ROM, etc., that may be included in some embodiments of computer system **1000** as system memory **1020** or another type of memory. In other embodiments, program instructions may be communicated using optical, acoustical or other form of propagated signal (e.g., carrier waves, infrared signals, digital signals, etc.) conveyed via a communication medium such as a network and/or a wireless link, such as may be implemented via network interface **1040**.

(105) It is noted that any of the distributed system embodiments described herein, or any of their components, may be implemented as one or more network-based services. For example, a compute cluster within a computing service may present computing services and/or other types of services that employ the distributed computing systems described herein to clients as network-based services. In some embodiments, a network-based service may be implemented by a software and/or hardware system designed to support interoperable machine-to-machine interaction over a network. A network-based service may have an interface described in a machine-processable format, such as the Web Services Description Language (WSDL). Other systems may interact with the network-based service in a manner prescribed by the description of the network-based service's interface. For example, the network-based service may define various operations that other systems may invoke and may define a particular application programming interface (API) to which other systems may be expected to conform when requesting the various operations.

(106) In various embodiments, a network-based service may be requested or invoked through the use of a message that includes parameters and/or data associated with the network-based services request. Such a message may be formatted according to a particular markup language such as Extensible Markup Language (XML), and/or may be encapsulated using a protocol such as Simple Object Access Protocol (SOAP). To perform a network-based services request, a network-based services client may assemble a message including the request and convey the message to an addressable endpoint (e.g., a Uniform Resource Locator (URL)) corresponding to the network-based service, using an Internet-based application layer transfer protocol such as Hypertext Transfer Protocol (HTTP).

(107) In some embodiments, network-based services may be implemented using Representational State Transfer (“RESTful”) techniques rather than message-based techniques. For example, a network-based service implemented according to a RESTful technique may be invoked through parameters included within an HTTP method such as PUT, GET, or DELETE, rather than encapsulated within a SOAP message.

(108) Although the embodiments above have been described in considerable detail, numerous variations and modifications may be made as would become apparent to those skilled in the art once the above disclosure is fully appreciated. It is intended that the following claims be interpreted to embrace all such modifications and changes and, accordingly, the above description to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A database system, comprising: a global shared memory; a one or more computing devices individually comprising at least one processor and a memory; a plurality of processes organized into a plurality of query processing groups, wherein individual ones of the query processing groups respectfully comprise two or more processes of the plurality of processes, wherein individual ones of the plurality of processes are configured to: allocate memory to a group-private shared memory arena from the global shared memory, wherein the group-private shared memory arena is allocated on behalf of a processing group from the plurality of query processing groups, the processing group comprising two or more processes of the plurality of processes; detect an abnormal execution condition for a failing process of the processing group during execution of a query, and responsive to the detecting: continue execution of other processes of the plurality of processes different from the two or more processes of the processing group while recovering execution of the processing group, wherein to recover execution of the processing group, the individual ones of the plurality of processes are configured to: restore execution of the two or more processes of the processing group to a normal execution state; and return the allocated memory of the group-private memory arena to the global shared memory.
2. The system of claim 1, wherein individual ones of the plurality of processing groups respectively comprise two or more processes having read-write access to group-private shared memory arenas and read-only access to a public shared memory arena.
3. The system of claim 1, wherein the failing process of the query processing group is a leader process, and wherein to restore execution of the two or more processes of the query processing group to the normal execution state, the individual ones of the plurality of processes are configured to: abort the query at the two or more processes of the query processing group; and restart the two or more processes of the processing group.
4. The system of claim 1, wherein the failing process of the query processing group is a compute process directed by a leader process of the query processing group, and wherein to restore execution of the two or more processes of the query processing group to the normal execution state, the individual ones of the plurality of processes are configured to: send a message, from the failing process to the leader process, indicating the detected abnormal execution condition; abort the query at the two or more processes of the query processing group; and restart the failing process.
5. A method, comprising: allocating memory to a group-private shared memory arena from a global shared memory of a computing system comprising a plurality of processes organized into a plurality of processing groups, wherein the group-private shared memory arena is allocated on behalf of a processing group from the plurality of processing groups, the processing group comprising two or more processes of the plurality of processes; detecting an abnormal execution condition for a failing process of the processing group during execution of a processing job, and responsive to the detecting: continuing execution of other processes of the plurality of processes different from the two or more processes of the processing group while recovering execution of the processing group, the recovering execution of the processing group comprising: restoring execution of the two or more processes of the processing group to a normal execution state; and returning the allocated memory of the group-private memory arena to the global shared memory.
6. The method of claim 5, wherein the plurality of processing groups respectively comprise processes having read-write access to group-private shared memory arenas and read-only access to a public shared memory arena.
7. The method of claim 6, further comprising: escalating privilege of a particular process of the processing group to obtain read-write access to the public shared memory arena; detecting an abnormal execution condition for the particular process while executing at the escalated privilege; and restarting the plurality of processes responsive to the detecting.
8. The method of claim 5, wherein the failing process of the processing group is a leader process, and wherein restoring execution of the two or more processes of the processing group to the normal

execution state comprises: aborting the processing job at the two or more processes of the processing group; and restarting the two or more processes of the processing group.

9. The method of claim 5, wherein the failing process of the processing group is a compute process directed by a leader process of the processing group, and wherein restoring execution of the two or more processes of the processing group to the normal execution state comprises: sending a message, from the failing process to the leader process, indicating the detected abnormal execution condition; aborting the processing job at the two or more processes of the processing group; and restarting the failing process.

10. The method of claim 5, the recovering execution of the processing group further comprising: signaling, to a dedicated thread of the failing process, information describing the abnormal execution condition.

11. The method of claim 10, further comprising: determining, by the dedicated thread of the failing process, a recovery state for the abnormal execution condition according to one or more execution privileges of the failing process; restarting the plurality of processes of the computing system responsive to determining that the recovery state is a safe recovery state; and causing, by the dedicated thread of the failing process, the returning of the allocated memory and the restarting of the two or more processes of the processing group responsive to determining that the recovery state is a safe recovery state.

12. The method of claim 5, wherein the computing system is a computing cluster comprising a plurality of computing nodes, and wherein the processing group is one of a plurality of processing groups respectively providing query processing for a distributed database service.

13. One or more non-transitory computer-accessible storage media storing program instructions that when executed on or across one or more computing devices cause the one or more computing devices to perform: allocating memory to a group-private shared memory arena from a global shared memory of a computing system comprising a plurality of processes organized into a plurality of query processing groups, wherein the group-private shared memory arena is allocated on behalf of a processing group from the plurality of processing groups, the processing group comprising two or more processes of the plurality of processes; detecting an abnormal execution condition for a failing process of the processing group during execution of a processing job, and responsive to the detecting: continuing execution of other processes of the plurality of processes different from the two or more processes of the processing group while recovering execution of the processing group, the recovering execution of the processing group comprising: restoring execution of the two or more processes of the processing group to a normal execution state; and returning the allocated memory of the group-private memory arena to the global shared memory.

14. The one or more non-transitory computer-accessible storage media of claim 13, wherein the processing group is one of a plurality of processing groups respectively comprising processes having read-write access to group-private shared memory arenas and read-only access to a public shared memory arena.

15. The one or more non-transitory computer-accessible storage media of claim 13, the program instructions that when executed on or across one or more computing devices cause the one or more computing devices to further perform: escalating privilege of a particular process of the processing group to obtain read-write access to the public shared memory area; detecting an abnormal execution condition for the particular process while executing at the escalated privilege; and restarting the plurality of processes responsive to the detecting.

16. The one or more non-transitory computer-accessible storage media of claim 13, wherein the failing process of the processing group is a leader process, and wherein restoring execution of the two or more processes of the processing group to the normal execution state comprises: canceling the processing job at the two or more processes of the processing group; and restarting the two or more processes of the processing group.

17. The one or more non-transitory computer-accessible storage media of claim 13, wherein the

failing process of the processing group is a compute process directed by a leader process of the processing group, and wherein restoring execution of the two or more processes of the processing group to the normal execution state comprises: sending a message, from the failing process to the leader process, indicating the detected abnormal execution condition; canceling the processing job at the two or more processes of the processing group; and restarting the failing process.

18. The one or more non-transitory computer-accessible storage media of claim 13, the recovering execution of the processing group further comprising: signaling, to a dedicated thread of the failing process, information describing the abnormal execution condition; and causing, by the dedicated thread of the failing process, the returning of the allocated memory and the restarting of the two or more processes of the processing group.

19. The one or more non-transitory computer-accessible storage media of claim 13, the program instructions that when executed on or across one or more computing devices cause the one or more computing devices to further perform: determining, by the dedicated thread of the failing process, a recovery state for the abnormal execution condition according to one or more execution privileges of the failing process; restarting the plurality of processes of the computing system responsive to determining that the recovery state is an safe recovery state; and causing, by the dedicated thread of the failing process, the returning of the allocated memory and the restarting of the two or more processes of the processing group responsive to determining that the recovery state is a safe recovery state.

20. The one or more non-transitory computer-accessible storage media of claim 13, wherein the computing system is a computing cluster comprising a plurality of computing nodes, and wherein the processing group is one of a plurality of processing groups respectively providing query processing for a distributed database service.
