

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386656
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Windh; Skyler Arron et al.

Dynamic decomposition and thread allocation

Abstract

Devices and techniques for thread scheduling control and memory splitting in a processor are described herein. An apparatus includes a hardware interface configured to receive a first request to execute a first thread, the first request including an indication of a workload; and processing circuitry configured to: determine the workload to produce a metric based at least in part on the indication; compare the metric with a threshold to determine that the metric is beyond the threshold; divide, based at least in part on the comparison, the workload into a set of sub-workloads consisting of predefined number of equal parts from the workload; create a second request to execute a second thread, the second request including a first member of the set of sub-workloads; and process a second member of the set of sub-workloads in the first thread.

Inventors: Windh; Skyler Arron (McKinney, TX), Brewer; Tony M. (Plano, TX), Estep; Patrick (Rowlett, TX)

Applicant: Micron Technology, Inc. (Boise, ID)

Family ID: 82119884

Assignee: Micron Technology, Inc. (Boise, ID)

Appl. No.: 17/465021

Filed: September 02, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220206846 A1	Jun. 30, 2022

Related U.S. Application Data

us-provisional-application US 63132754 20201231

Publication Classification

Int. Cl.: **G06F9/48** (20060101); **G06F9/38** (20180101); **G06F9/50** (20060101); **G06F11/34** (20060101)

U.S. Cl.:

CPC **G06F9/4881** (20130101); **G06F9/3851** (20130101); **G06F9/3888** (20230801); **G06F9/5038** (20130101); **G06F11/3409** (20130101); G06F2209/501 (20130101); G06F2209/5022 (20130101)

Field of Classification Search

CPC: G06F (9/4881); G06F (9/5038); G06F (11/3409); G06F (2209/501); G06F (2209/5022)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5434972	12/1994	Hamlin	370/408	G06F 9/4494
8169437	12/2011	Legakis	345/428	G06T 17/20
8229946	12/2011	Crane	712/28	G06F 9/505
8813091	12/2013	Maessen	718/105	G06F 9/5066
9043401	12/2014	Wong	709/205	G06F 9/505
2007/0061519	12/2006	Barrett et al.	N/A	N/A
2007/0067606	12/2006	Lin	712/10	G06F 9/5044
2010/0251257	12/2009	Kim	718/105	G06F 9/5027
2012/0246158	12/2011	Ke	707/E17.001	G06F 16/278
2014/0244643	12/2013	Basak	707/737	G06F 11/3414
2015/0026698	12/2014	Malakhov	718/105	G06F 9/5077
2015/0103677	12/2014	Lee	370/249	H04L 41/0896
2015/0193959	12/2014	Shah	345/473	G06F 3/00
2016/0011901	12/2015	Hurwitz	718/101	G06F 9/46
2016/0179574	12/2015	Merrill, III	718/106	G06F 7/523
2016/0335127	12/2015	Artmeier	N/A	G06F 9/5044
2018/0103088	12/2017	Blainey	N/A	H04L 67/1008
2019/0303387	12/2018	Smarda	N/A	G06F 9/4843
2020/0184366	12/2019	Mandal	N/A	G06N 3/042
2020/0233706	12/2019	Smith	N/A	G06F 9/5066
2020/0280511	12/2019	Gapin et al.	N/A	N/A
2021/0019219	12/2020	Patel	N/A	G06F 9/5066
2021/0124611	12/2020	Saillet	N/A	G06F 9/4887
2021/0183005	12/2020	Du	N/A	G06F 9/30032
2021/0342200	12/2020	Gupta	N/A	G06F 1/08
2023/0031226	12/2022	Lee	N/A	G06N 3/063

OTHER PUBLICATIONS

“Cilk”, [Online]. Retrieved from the Internet: URL: <https://web.archive.org/web/20200412001219/https://en.wikipedia.org/wiki/Cilk>, (Accessed on Apr. 12, 2020), 10 pgs. cited by applicant

“Chinese Application Serial No. 202111626756.3, Office Action mailed Nov. 26, 2024”, w/ English translation, 18 pgs. cited by applicant
“Chinese Application Serial No. 202111626756.3, Response filed Mar. 20, 2025 to Office Action mailed Nov. 26, 2024”, with English claims, 16 pages, English claims only. cited by applicant

Primary Examiner: Wu; Benjamin C

Attorney, Agent or Firm: Schwegman Lundberg & Woessner, P.A.

Background/Summary

RELATED APPLICATION (1) This application claims priority to U.S. Provisional Application Ser. No. 63/132,754, filed Dec. 31, 2020, which is incorporated herein by reference in its entirety.

BACKGROUND

(1) Various computer architectures, such as the Von Neumann architecture, conventionally use a shared memory for data, a bus for accessing the shared memory, an arithmetic unit, and a program control unit. However, moving data between processors and memory can require significant time and energy, which in turn can constrain performance and capacity of computer systems. In view of these limitations, new computing architectures and devices are desired to advance computing performance beyond the practice of transistor scaling (i.e., Moore's Law).

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The disclosure will be understood more fully from the detailed description given below and from the accompanying drawings of various embodiments of the disclosure. The drawings, however, should not be taken to limit the disclosure to the specific embodiments, but are for explanation and understanding only.

(2) To easily identify the discussion of any particular element or act, the most significant digit or digits in a reference number refer to the figure number in which that element is first introduced.

(3) FIG. 1 illustrates generally a first example of a first memory-compute device in the context of a memory-compute system, according to an embodiment.

(4) FIG. 2 illustrates generally an example of a memory subsystem of a memory-compute device, according to an embodiment.

(5) FIG. 3 illustrates generally an example of a programmable atomic unit for a memory controller, according to an embodiment.

(6) FIG. 4 illustrates an example of a hybrid threading processor (HTP) accelerator of a memory-compute device, according to an embodiment.

(7) FIG. 5 illustrates an example of a representation of a hybrid threading fabric (HTF) of a memory-compute device, according to an embodiment.

(8) FIG. 6A illustrates generally an example of a chiplet system, according to an embodiment.

(9) FIG. 6B illustrates generally a block diagram showing various components in the chiplet system from the example of FIG. 6A.

(10) FIG. 7 illustrates generally an example of a chiplet-based implementation for a memory-compute device, according to an embodiment.

(11) FIG. 8 illustrates an example tiling of memory-compute device chiplets, according to an embodiment.

(12) FIG. 9 is a table that illustrates fiber create instructions, according to an embodiment.

(13) FIG. 10 is a pseudocode listing illustrating a divide and conquer process, according to an embodiment.

(14) FIG. 11 is another code listing illustrating a divide and conquer process, according to an embodiment.

(15) FIG. 12 is a diagram illustrating a thread calling tree, according to an embodiment.

(16) FIG. 13 is a flow chart of an example of a method for thread execution control, according to an embodiment.

(17) FIG. 14 illustrates a block diagram of an example machine with which, in which, or by which any one or more of the techniques (e.g., methodologies) discussed herein can be implemented.

DETAILED DESCRIPTION

(18) Recent advances in materials, devices, and integration technology, can be leveraged to provide memory-centric compute topologies. Such topologies can realize advances in compute efficiency and workload throughput, for example, for applications constrained by size, weight, or power requirements. The topologies can be used to facilitate low-latency compute near, or inside of, memory or other data storage elements. The approaches can be particularly well-suited for various compute-intensive operations with sparse lookups, such as in transform computations (e.g., fast Fourier transform computations (FFT)), or in applications such as neural networks or artificial intelligence (AI), financial analytics, or simulations or modeling such as for computational fluid dynamics (CFD), Enhanced Acoustic Simulator for Engineers (EASE), Simulation Program with Integrated Circuit Emphasis (SPICE), and others.

(19) Systems, devices, and methods discussed herein can include or use memory-compute systems with processors, or processing capabilities, that are provided in, near, or integrated with memory or data storage components. Such systems are referred to generally herein as compute-near-memory (CNM) systems. A CNM system can be a node-based system with individual nodes in the systems coupled using a system scale fabric. Each node can include or use specialized or general purpose processors, and user-accessible accelerators, with a custom compute fabric to facilitate intensive operations, particularly in environments where high cache miss rates are expected.

(20) In an example, each node in a CNM system can have a host processor or processors. Within each node, a dedicated hybrid threading processor can occupy a discrete endpoint of an on-chip network. The hybrid threading processor can have access to some or all of the memory in a particular node of the system, or a hybrid threading processor can have access to memories across a network of multiple nodes via the system scale fabric. The custom compute fabric, or hybrid threading fabric, at each node can have its own processor(s) or accelerator(s) and can operate at higher bandwidth than the hybrid threading processor. Different nodes in a compute-near-memory system can be differently configured, such as having different compute capabilities, different types of memories, different interfaces, or other differences. However, the nodes can be commonly coupled to share data and compute resources within a defined address space.

(21) In an example, a compute-near-memory system, or a node within the system, can be user-configured for custom operations. A user can provide instructions using a high-level programming language, such as C/C++, that can be compiled and mapped directly into a dataflow architecture of the system, or of one or more nodes in the CNM system. That is, the nodes in the system can include hardware blocks (e.g., memory controllers, atomic units, other customer accelerators, etc.) that can be configured to directly implement or support user instructions to thereby enhance system performance and reduce latency.

(22) In an example, a compute-near-memory system can be particularly suited for implementing a hierarchy of instructions and nested loops (e.g., two, three, or more, loops deep, or multiple-dimensional loops). A standard compiler can be used to accept high-level language instructions and, in turn, compile directly into the dataflow architecture of one or more of the nodes. For example, a node in the system can include a hybrid threading fabric accelerator. The hybrid

threading fabric accelerator can execute in a user space of the CNM system and can initiate its own threads or sub-threads, which can operate in parallel. Each thread can map to a different loop iteration to thereby support multi-dimensional loops. With the capability to initiate such nested loops, among other capabilities, the CNM system can realize significant time savings and latency improvements for compute-intensive operations.

(23) A compute-near-memory system, or nodes or components of a compute-near-memory system, can include or use various memory devices, controllers, and interconnects, among other things. In an example, the system can comprise various interconnected nodes and the nodes, or groups of nodes, can be implemented using chiplets. Chiplets are an emerging technique for integrating various processing functionality. Generally, a chiplet system is made up of discrete chips (e.g., integrated circuits (ICs) on different substrate or die) that are integrated on an interposer and packaged together. This arrangement is distinct from single chips (e.g., ICs) that contain distinct device blocks (e.g., intellectual property (IP) blocks) on one substrate (e.g., single die), such as a system-on-a-chip (SoC), or discretely packaged devices integrated on a board. In general, chiplets provide production benefits than single die chips, including higher yields or reduced development costs. FIG. 6A and FIG. 6B, discussed below, illustrate generally an example of a chiplet system such as can comprise a compute-near-memory system.

(24) In some computing tasks, hundreds or even thousands of threads may be used for parallel processing. At this scale, such as with thousands of threads, the amount of work able to be parallelized is a key factor in achieving performance. Another consideration is how quickly the threads can be started.

(25) To process large lists in parallel, one strategy is to have a master thread that works on dividing the list into chunks and creating child threads to process the chunks. This can be done with very little scheduling overhead, and the loop can iterate quickly. However, this strategy has two main drawbacks. First, if there are many chunks to start, a single source of work generation means the time it takes to get every unit of work started becomes significant. Second, this linear scheduling of tasks also creates pressure in the tasking interfaces as there is a single point of joining and freeing child threads.

(26) What is needed is a more efficient mechanism to process large tasks in parallel. The systems and methods described here partition the work using a divide and conquer strategy. Instead of a single master thread creating all of the work threads, every thread participates in the scheduling. If the amount of work assigned to a thread is larger than some threshold size chunk, the thread first attempts to split the work in half. The thread creates a new thread to process one half of the work and continues working with the remaining half. If the thread creation fails (e.g., all system threads are already running), then the thread will process one threshold-sized chunk of work before trying to split again. This mechanism provides a logarithmic schedule tree of threads.

(27) Efficiency can be further improved by taking advantage of a busy fail thread creation application programming interface (API) that the hybrid threading processor (HTP) runtime provides. A logarithmic thread creation tree also as a log deep return path. If a single result is needed from each sub task, then this must be done. However, if no result is needed (e.g., all results are already in memory), then we can take advantage of a special “no return” variant of thread creation. Threads without a return are automatically terminated by the host interface as soon as they finish executing. That means that they are again free to be reused, and the calling code does not have to handle joining and freeing the thread. This greatly increases the flexibility of the work splitting, as threads can easily be continually recycled as new work is generated.

(28) Using this strategy, work continues to be processed and dynamic thread creation does not deadlock. An additional advantage is that the scheduling mechanism results in a balanced tree schedule of logarithmic depth with an exponential number of threads running in few levels. Further processing advantages may be gained though the use of “Busy Fail” threads that provide quick dispatch and return to all resources in a system and “no return” thread variants that enable quick

termination and reuse of system threads.

(29) A compute-near-memory system, or nodes or components of a compute-near-memory system, can include or use various memory devices, controllers, and interconnect fabrics, among other things. In an example, the system can comprise a computing fabric with various interconnected nodes and the nodes, or groups of nodes, can be implemented using chiplets. Chiplets are an emerging technique for integrating various processing functionality. Generally, a chiplet system is made up of discrete chips (e.g., integrated circuits (ICs) on different substrate or die) that are integrated on an interposer and packaged together. This arrangement is distinct from single chips (e.g., ICs) that contain distinct device blocks (e.g., intellectual property (IP) blocks) on one substrate (e.g., single die), such as a system-on-a-chip (SoC), or discretely packaged devices integrated on a board. In general, chiplets provide better performance (e.g., lower power consumption, reduced latency, etc.) than discretely packaged devices, and chiplets provide greater production benefits than single die chips. These production benefits can include higher yields or reduced development costs and time.

(30) FIG. 1 illustrates generally a first example of a compute-near-memory system, or CNM system **102**. The example of the CNM system **102** includes multiple different memory-compute nodes, such as can each include various compute-near-memory devices. Each node in the system can operate in its own operating system (OS) domain (e.g., Linux, among others). In an example, the nodes can exist collectively in a common OS domain of the CNM system **102**.

(31) The example of FIG. 1 includes an example of a first memory-compute node **104** of the CNM system **102**. The CNM system **102** can have multiple nodes, such as including different instances of the first memory-compute node **104**, that are coupled using a scale fabric **106**. In an example, the architecture of the CNM system **102** can support scaling with up to a different memory-compute nodes (e.g., n=4096) using the scale fabric **106**. As further discussed below, each node in the CNM system **102** can be an assembly of multiple devices.

(32) The CNM system **102** can include a global controller for the various nodes in the system, or a particular memory-compute node in the system can optionally serve as a host or controller to one or multiple other memory-compute nodes in the same system. The various nodes in the CNM system **102** can thus be similarly or differently configured.

(33) In an example, each node in the CNM system **102** can comprise a host system that uses a specified operating system. The operating system can be common or different among the various nodes in the CNM system **102**. In the example of FIG. 1, the first memory-compute node **104** comprises a host system **108**, a first switch **110**, and a first memory-compute device **112**. The host system **108** can comprise a processor, such as can include an X86, ARM, RISC-V, or other type of processor. The first switch **110** can be configured to facilitate communication between or among devices of the first memory-compute node **104** or of the CNM system **102**, such as using a specialized or other communication protocol, generally referred to herein as a chip-to-chip protocol interface (CTCPI). That is, the CTCPI can include a specialized interface that is unique to the CNM system **102**, or can include or use other interfaces such as the compute express link (CXL) interface, the peripheral component interconnect express (PCIe) interface, or the chiplet protocol interface (CPI), among others. The first switch **110** can include a switch configured to use the CTCPI. For example, the first switch **110** can include a CXL switch, a PCIe switch, a CPI switch, or other type of switch. In an example, the first switch **110** can be configured to couple differently configured endpoints. For example, the first switch **110** can be configured to convert packet formats, such as between PCIe and CPI formats, among others.

(34) The CNM system **102** is described herein in various example configurations, such as comprising a system of nodes, and each node can comprise various chips (e.g., a processor, a switch, a memory device, etc.). In an example, the first memory-compute node **104** in the CNM system **102** can include various chips implemented using chiplets. In the below-discussed chiplet-based configuration of the CNM system **102**, inter-chiplet communications, as well as additional

communications within the system, can use a CPI network. The CPI network described herein is an example of the CTCPI, that is, as a chiplet-specific implementation of the CTCPI. As a result, the below-described structure, operations, and functionality of CPI can apply equally to structures, operations, and functions as may be otherwise implemented using non-chiplet-based CTCPI implementations. Unless expressly indicated otherwise, any discussion herein of CPI applies equally to CTCPI.

(35) A CPI interface includes a packet-based network that supports virtual channels to enable a flexible and high-speed interaction between chiplets, such as can comprise portions of the first memory-compute node **104** or the CNM system **102**. The CPI can enable bridging from intra-chiplet networks to a broader chiplet network. For example, the Advanced eXtensible Interface (AXI) is a specification for intra-chip communications. AXI specifications, however, cover a variety of physical design options, such as the number of physical channels, signal timing, power, etc. Within a single chip, these options are generally selected to meet design goals, such as power consumption, speed, etc. However, to achieve the flexibility of a chiplet-based memory-compute system, an adapter, such as using CPI, can interface between the various AXI design options that can be implemented in the various chiplets. By enabling a physical channel-to-virtual channel mapping and encapsulating time-based signaling with a packetized protocol, CPI can be used to bridge intra-chiplet networks, such as within a particular memory-compute node, across a broader chiplet network, such as across the first memory-compute node **104** or across the CNM system **102**.

(36) The CNM system **102** is scalable to include multiple-node configurations. That is, multiple different instances of the first memory-compute node **104**, or of other differently configured memory-compute nodes, can be coupled using the scale fabric **106**, to provide a scaled system. Each of the memory-compute nodes can run its own operating system and can be configured to jointly coordinate system-wide resource usage.

(37) In the example of FIG. **1**, the first switch **110** of the first memory-compute node **104** is coupled to the scale fabric **106**. The scale fabric **106** can provide a switch (e.g., a CTCPI switch, a PCIe switch, a CPI switch, or other switch) that can facilitate communication among and between different memory-compute nodes. In an example, the scale fabric **106** can help various nodes communicate in a partitioned global address space (PGAS).

(38) In an example, the first switch **110** from the first memory-compute node **104** is coupled to one or multiple different memory-compute devices, such as including the first memory-compute device **112**. The first memory-compute device **112** can comprise a chiplet-based architecture referred to herein as a compute-near-memory (CNM) chiplet. A packaged version of the first memory-compute device **112** can include, for example, one or multiple CNM chiplets. The chiplets can be communicatively coupled using CTCPI for high bandwidth and low latency.

(39) In the example of FIG. **1**, the first memory-compute device **112** can include a network on chip (NOC) or first NOC **118**. Generally, a NOC is an interconnection network within a device, connecting a particular set of endpoints. In FIG. **1**, the first NOC **118** can provide communications and connectivity between the various memory, compute resources, and ports of the first memory-compute device **112**.

(40) In an example, the first NOC **118** can comprise a folded Clos topology, such as within each instance of a memory-compute device, or as a mesh that couples multiple memory-compute devices in a node. The Clos topology, such as can use multiple, smaller radix crossbars to provide functionality associated with a higher radix crossbar topology, offers various benefits. For example, the Clos topology can exhibit consistent latency and bisection bandwidth across the NOC.

(41) The first NOC **118** can include various distinct switch types including hub switches, edge switches, and endpoint switches. Each of the switches can be constructed as crossbars that provide substantially uniform latency and bandwidth between input and output nodes. In an example, the endpoint switches and the edge switches can include two separate crossbars, one for traffic headed to the hub switches, and the other for traffic headed away from the hub switches. The hub switches

can be constructed as a single crossbar that switches all inputs to all outputs.

(42) In an example, the hub switches can have multiple ports each (e.g., four or six ports each), such as depending on whether the particular hub switch participates in inter-chip communications. A number of hub switches that participates in inter-chip communications can be set by an inter-chip bandwidth requirement.

(43) The first NOC **118** can support various payloads (e.g., from 8 to 64-byte payloads; other payload sizes can similarly be used) between compute elements and memory. In an example, the first NOC **118** can be optimized for relatively smaller payloads (e.g., 8-16 bytes) to efficiently handle access to sparse data structures.

(44) In an example, the first NOC **118** can be coupled to an external host via a first physical-layer interface **114**, a PCIe subordinate module **116** or endpoint, and a PCIe principal module **126** or root port. That is, the first physical-layer interface **114** can include an interface to allow an external host processor to be coupled to the first memory-compute device **112**. An external host processor can optionally be coupled to one or multiple different memory-compute devices, such as using a PCIe switch or other, native protocol switch. Communication with the external host processor through a PCIe-based switch can limit device-to-device communication to that supported by the switch. Communication through a memory-compute device-native protocol switch such as using CTCPI, in contrast, can allow for more full communication between or among different memory-compute devices, including support for a partitioned global address space, such as for creating threads of work and sending events.

(45) In an example, the CTCPI protocol can be used by the first NOC **118** in the first memory-compute device **112**, and the first switch **110** can include a CTCPI switch. The CTCPI switch can allow CTCPI packets to be transferred from a source memory-compute device, such as the first memory-compute device **112**, to a different, destination memory-compute device (e.g., on the same or other node), such as without being converted to another packet format.

(46) In an example, the first memory-compute device **112** can include an internal host processor **122**. The internal host processor **122** can be configured to communicate with the first NOC **118** or other components or modules of the first memory-compute device **112**, for example, using the internal PCIe principal module **126**, which can help eliminate a physical layer that would consume time and energy. In an example, the internal host processor **122** can be based on a RISC-V ISA processor, and can use the first physical-layer interface **114** to communicate outside of the first memory-compute device **112**, such as to other storage, networking, or other peripherals to the first memory-compute device **112**. The internal host processor **122** can control the first memory-compute device **112** and can act as a proxy for operating system-related functionality. The internal host processor **122** can include a relatively small number of processing cores (e.g., 2-4 cores) and a host memory device **124** (e.g., comprising a DRAM module).

(47) In an example, the internal host processor **122** can include PCI root ports. When the internal host processor **122** is in use, then one of its root ports can be connected to the PCIe subordinate module **116**. Another of the root ports of the internal host processor **122** can be connected to the first physical-layer interface **114**, such as to provide communication with external PCI peripherals. When the internal host processor **122** is disabled, then the PCIe subordinate module **116** can be coupled to the first physical-layer interface **114** to allow an external host processor to communicate with the first NOC **118**. In an example of a system with multiple memory-compute devices, the first memory-compute device **112** can be configured to act as a system host or controller. In this example, the internal host processor **122** can be in use, and other instances of internal host processors in the respective other memory-compute devices can be disabled.

(48) The internal host processor **122** can be configured at power-up of the first memory-compute device **112**, such as to allow the host to initialize. In an example, the internal host processor **122** and its associated data paths (e.g., including the first physical-layer interface **114**, the PCIe subordinate module **116**, etc.) can be configured from input pins to the first memory-compute

device **112**. One or more of the pins can be used to enable or disable the internal host processor **122** and configure the PCI (or other) data paths accordingly.

(49) In an example, the first NOC **118** can be coupled to the scale fabric **106** via a scale fabric interface module **136** and a second physical-layer interface **138**. The scale fabric interface module **136**, or SIF, can facilitate communication between the first memory-compute device **112** and a device space, such as a partitioned global address space (PGAS). The PGAS can be configured such that a particular memory-compute device, such as the first memory-compute device **112**, can access memory or other resources on a different memory-compute device (e.g., on the same or different node), such as using a load/store paradigm. Various scalable fabric technologies can be used, including CTCPI, CPI, Gen-Z, PCI, or Ethernet bridged over CXL. The scale fabric **106** can be configured to support various packet formats. In an example, the scale fabric **106** supports orderless packet communications, or supports ordered packets such as can use a path identifier to spread bandwidth across multiple equivalent paths. The scale fabric **106** can generally support remote operations such as remote memory read, write, and other built-in atomics, remote memory atomics, remote memory-compute device send events, and remote memory-compute device call and return operations.

(50) In an example, the first NOC **118** can be coupled to one or multiple different memory modules, such as including a first memory device **128**. The first memory device **128** can include various kinds of memory devices, for example, LPDDR5 or GDDR6, among others. In the example of FIG. **1**, the first NOC **118** can coordinate communications with the first memory device **128** via a memory controller **130** that can be dedicated to the particular memory module. In an example, the memory controller **130** can include a memory module cache and an atomic operations module. The atomic operations module can be configured to provide relatively high-throughput atomic operators, such as including integer and floating-point operators. The atomic operations module can be configured to apply its operators to data within the memory module cache (e.g., comprising SRAM memory side cache), thereby allowing back-to-back atomic operations using the same memory location, with minimal throughput degradation.

(51) The memory module cache can provide storage for frequently accessed memory locations, such as without having to re-access the first memory device **128**. In an example, the memory module cache can be configured to cache data only for a particular instance of the memory controller **130**. In an example, the memory controller **130** includes a DRAM controller configured to interface with the first memory device **128**, such as including DRAM devices. The memory controller **130** can provide access scheduling and bit error management, among other functions.

(52) In an example, the first NOC **118** can be coupled to a hybrid threading processor (HTP **140**), a hybrid threading fabric (HTF **142**) and a host interface and dispatch module (HIF **120**). The HIF **120** can be configured to facilitate access to host-based command request queues and response queues. In an example, the HIF **120** can dispatch new threads of execution on processor or compute elements of the HTP **140** or the HTF **142**. In an example, the HIF **120** can be configured to maintain workload balance across the HTP **140** module and the HTF **142** module.

(53) The hybrid threading processor, or HTP **140**, can include an accelerator, such as can be based on a RISC-V instruction set. The HTP **140** can include a highly threaded, event-driven processor in which threads can be executed in single instruction rotation, such as to maintain high instruction throughput. The HTP **140** comprises relatively few custom instructions to support low-overhead threading capabilities, event send/receive, and shared memory atomic operators.

(54) The hybrid threading fabric, or HTF **142**, can include an accelerator, such as can include a non-von Neumann, coarse-grained, reconfigurable processor. The HTF **142** can be optimized for high-level language operations and data types (e.g., integer or floating point). In an example, the HTF **142** can support data flow computing. The HTF **142** can be configured to use substantially all of the memory bandwidth available on the first memory-compute device **112**, such as when executing memory-bound compute kernels.

(55) The HTP and HTF accelerators of the CNM system **102** can be programmed using various high-level, structured programming languages. For example, the HTP and HTF accelerators can be programmed using C/C++, such as using the LLVM compiler framework. The HTP accelerator can leverage an open source compiler environment, such as with various added custom instruction sets configured to improve memory access efficiency, provide a message passing mechanism, and manage events, among other things. In an example, the HTF accelerator can be designed to enable programming of the HTF **142** using a high-level programming language, and the compiler can generate a simulator configuration file or a binary file that runs on the HTF **142** hardware. The HTF **142** can provide a mid-level language for expressing algorithms precisely and concisely, while hiding configuration details of the HTF accelerator itself. In an example, the HTF accelerator tool chain can use an LLVM front-end compiler and the LLVM intermediate representation (IR) to interface with an HTF accelerator back end.

(56) FIG. 2 illustrates generally an example of a memory subsystem **200** of a memory-compute device, according to an embodiment. The example of the memory subsystem **200** includes a controller **202**, a programmable atomic unit **208**, and a second NOC **206**. The controller **202** can include or use the programmable atomic unit **208** to carry out operations using information in a memory device **204**. In an example, the memory subsystem **200** comprises a portion of the first memory-compute device **112** from the example of FIG. 1, such as including portions of the first NOC **118** or of the memory controller **130**.

(57) In the example of FIG. 2, the second NOC **206** is coupled to the controller **202** and the controller **202** can include a memory control module **210**, a local cache module **212**, and a built-in atomics module **214**. In an example, the built-in atomics module **214** can be configured to handle relatively simple, single-cycle, integer atomics. The built-in atomics module **214** can perform atomics at the same throughput as, for example, normal memory read or write operations. In an example, an atomic memory operation can include a combination of storing data to the memory, performing an atomic memory operation, and then responding with load data from the memory.

(58) The local cache module **212**, such as can include an SRAM cache, can be provided to help reduce latency for repetitively-accessed memory locations. In an example, the local cache module **212** can provide a read buffer for sub-memory line accesses. The local cache module **212** can be particularly beneficial for compute elements that have relatively small or no data caches.

(59) The memory control module **210**, such as can include a DRAM controller, can provide low-level request buffering and scheduling, such as to provide efficient access to the memory device **204**, such as can include a DRAM device. In an example, the memory device **204** can include or use a GDDR6 DRAM device, such as having 16 Gb density and 64 Gb/sec peak bandwidth. Other devices can similarly be used.

(60) In an example, the programmable atomic unit **208** can comprise single-cycle or multiple-cycle operator such as can be configured to perform integer addition or more complicated multiple-instruction operations such as bloom filter insert. In an example, the programmable atomic unit **208** can be configured to perform load and store-to-memory operations. The programmable atomic unit **208** can be configured to leverage the RISC-V ISA with a set of specialized instructions to facilitate interactions with the controller **202** to atomically perform user-defined operations.

(61) Programmable atomic requests, such as received from an on-node or off-node host, can be routed to the programmable atomic unit **208** via the second NOC **206** and the controller **202**. In an example, custom atomic operations (e.g., carried out by the programmable atomic unit **208**) can be identical to built-in atomic operations (e.g., carried out by the built-in atomics module **214**) except that a programmable atomic operation can be defined or programmed by the user rather than the system architect. In an example, programmable atomic request packets can be sent through the second NOC **206** to the controller **202**, and the controller **202** can identify the request as a custom atomic. The controller **202** can then forward the identified request to the programmable atomic unit **208**.

(62) FIG. 3 illustrates generally an example of a programmable atomic unit **302** for use with a memory controller, according to an embodiment. In an example, the programmable atomic unit **302** can comprise or correspond to the programmable atomic unit **208** from the example of FIG. 2. That is, FIG. 3 illustrates components in an example of a programmable atomic unit **302** (PAU), such as those noted above with respect to FIG. 2 (e.g., in the programmable atomic unit **208**), or to FIG. 1 (e.g., in an atomic operations module of the memory controller **130**). As illustrated in FIG. 3, the programmable atomic unit **302** includes a PAU processor or PAU core **306**, a PAU thread control **304**, an instruction SRAM **308**, a data cache **310** and a memory interface **312** to interface with the memory controller **314**. In an example, the memory controller **314** comprises an example of the controller **202** from the example of FIG. 2.

(63) In an example, the PAU core **306** is a pipelined processor such that multiple stages of different instructions are executed together per clock cycle. The PAU core **306** can include a barrel-multithreaded processor, with thread control **304** circuitry to switch between different register files (e.g., sets of registers containing current processing state) upon each clock cycle. This enables efficient context switching between currently executing threads. In an example, the PAU core **306** supports eight threads, resulting in eight register files. In an example, some or all of the register files are not integrated into the PAU core **306**, but rather reside in a local data cache **310** or the instruction SRAM **308**. This reduces circuit complexity in the PAU core **306** by eliminating the traditional flip-flops used for registers in such memories.

(64) The local PAU memory can include instruction SRAM **308**, such as can include instructions for various atomics. The instructions comprise sets of instructions to support various application-loaded atomic operators. When an atomic operator is requested, such as by an application chiplet, a set of instructions corresponding to the atomic operator are executed by the PAU core **306**. In an example, the instruction SRAM **308** can be partitioned to establish the sets of instructions. In this example, the specific programmable atomic operator being requested by a requesting process can identify the programmable atomic operator by the partition number. The partition number can be established when the programmable atomic operator is registered with (e.g., loaded onto) the programmable atomic unit **302**. Other metadata for the programmable instructions can be stored in memory (e.g., in partition tables) in memory local to the programmable atomic unit **302**.

(65) In an example, atomic operators manipulate the instruction SRAM **308**, which is generally synchronized (e.g., flushed) when a thread for an atomic operator completes. Thus, aside from initial loading from the external memory, such as from the memory controller **314**, latency can be reduced for most memory operations during execution of a programmable atomic operator thread.

(66) A pipelined processor, such as the PAU core **306**, can experience an issue when an executing thread attempts to issue a memory request if an underlying hazard condition would prevent such a request. Here, the memory request is to retrieve data from the memory controller **314**, whether it be from a cache on the memory controller **314** or off-die memory. To resolve this issue, the PAU core **306** is configured to deny the memory request for a thread. Generally, the PAU core **306** or the thread control **304** can include circuitry to enable one or more thread rescheduling points in the pipeline. Here, the denial occurs at a point in the pipeline that is beyond (e.g., after) these thread rescheduling points. In an example, the hazard occurred beyond the rescheduling point. Here, a preceding instruction in the thread created the hazard after the memory request instruction passed the last thread rescheduling point prior to the pipeline stage in which the memory request could be made.

(67) In an example, to deny the memory request, the U core **306** is configured to determine detect) that there is a hazard on memory indicated in the memory request. Here, hazard denotes any condition such that allowing (e.g., performing) the memory request will result in an inconsistent state for the thread. In an example, the hazard is an in-flight memory request. Here, whether or not the instruction SRAM **308** includes data for the requested memory address, the presence of the in-flight memory request, makes it uncertain what the data in the instruction SRAM **308** at that

address should be. Thus, the thread must wait for the in-flight memory request to be completed to operate on current data. The hazard is cleared when the memory request completes.

(68) In an example, the hazard is a dirty cache line in the instruction SRAM **308** for the requested memory address. Although the dirty cache line generally indicates that the data in the cache is current and the memory controller version of this data is not, an issue can arise on thread instructions that do not operate from the cache. An example of such an instruction uses a built-in atomic operator, or other separate hardware block, of the memory controller **314**. In the context of a memory controller, the built-in atomic operators can be separate from the programmable atomic unit **302** and do not have access to the cache or instruction SRAM **308** inside the PAU. Thus, a memory request for a built-in atomic operator is processed from the cache or the off-die memory. If the cache line is dirty, then the built-in atomic operator will not be operating on the most current data until the cache is flushed to synchronize the cache and the other or off-die memories. This same situation could occur with other hardware blocks of the memory controller, such as cryptography block, encoder, etc.

(69) The PAU core **306** can be configured to place a thread into a memory response path. Although the memory request of a thread was not issued, the thread can be largely handled as if the memory request was made. This enables a smooth rescheduling of the thread without impacting other threads and without adding thread stalling circuitry to the various pipeline stages of the PAU core **306**. Two issues can arise with this approach. First, because a memory request was not actually issued, the thread can skip some stages of the memory request and response path. Second, because no memory response is coming, data will not be updated (e.g., there will be no register writeback) by the thread.

(70) To address the first issue, the PAU core **306** can be configured to add the thread to a queue of threads awaiting insertion into the memory response path. Here, the PAU core **306** maintains the queue and pushes the thread onto the queue. In this way, the memory response path of the processor core can pop a next thread off the queue to place into the standard memory response path.

(71) To address the second issue, the PAU core **306** is configured to provide an indicator that a register write-back will not occur for the thread. The indicator (e.g., flag) can be some bits added to the thread metadata in the queue, or elsewhere, to indicate that a write-back mechanism will not operate with respect to the thread in this instance. Thus, when the thread appears at a register file write block in the memory response path, no register file write will occur.

(72) In an example, to provide the indicator that a register write-back will not occur for the thread, the PAU core **306** is configured to create a memory response that includes the indicator. Here, the way in which the thread is inserted into the memory response path is through a memory response manufactured by the PAU core **306** rather than from the memory controller **314**. Thus, to the memory response path, a standard memory response is being handled for the thread. However, the inclusion of the indicator means that a register or cache write block will not perform any actions. In an example, the memory response is placed into the queue to await insertion into the memory response path after earlier stages in the memory response path, such as decode from the memory controller, error correction, etc.

(73) The thread control **304** can be configured to reschedule a thread with other threads in a memory response path. In an example, the thread is popped from the queue of threads awaiting insertion into the memory response path. In an example, the indicator is processed by the memory response path to move the thread from the queue of threads awaiting a response from memory to a memory response reschedule queue in the memory response path. Here, the indicator, or other metadata, is used by the memory response path to determine a correct insertion time for the thread amongst other pending threads.

(74) FIG. 4 illustrates an example of a hybrid threading processor (HTP) accelerator, or HTP accelerator **400**. The HTP accelerator **400** can comprise a portion of a memory-compute device,

according to an embodiment. In an example, the HTP accelerator **400** can include or comprise the HTP **140** from the example of FIG. **1**. The HTP accelerator **400** includes, for example, a HTP core **402**, an instruction cache **404**, a data cache **406**, a translation block **408**, a memory interface **410**, and a thread controller **412**. The HTP accelerator **400** can further include a dispatch interface **414** and a NOC interface **416**, such as for interfacing with a NOC such as the first NOC **118** from the example of FIG. **1**, the second NOC **206** from the example of FIG. **2**, or other NOC.

(75) In an example, the HTP accelerator **400** includes a module that is based on a RISC-V instruction set, and can include a relatively small number of other or additional custom instructions to support a low-overhead, threading-capable Hybrid Threading (HT) language. The HTP accelerator **400** can include a highly-threaded processor core, the HTP core **402**, in which, or with which, threads can be executed in a single instruction rotation, such as to maintain high instruction throughput. In an example, a thread can be paused when it waits for other, pending events to complete. This can allow the compute resources to be efficiently used on relevant work instead of polling. In an example, multiple-thread barrier synchronization can use efficient HTP-to-HTP and HTP-to/from-Host messaging, such as can allow thousands of threads to initialize or wake in, for example, tens of clock cycles.

(76) In an example, the dispatch interface **414** can comprise a functional block of the HTP accelerator **400** for handling hardware-based thread management. That is, the dispatch interface **414** can manage dispatch of work to the HTP core **402** or other accelerators. Non-HTP accelerators, however, are generally not able to dispatch work. In an example, work dispatched from a host can use dispatch queues that reside in, e.g., host main memory (e.g., DRAM-based memory). Work dispatched from the HTP accelerator **400**, on the other hand, can use dispatch queues that reside in SRAM, such as within the dispatches for the target HTP accelerator **400** within a particular node.

(77) In an example, the HTP core **402** can comprise one or more cores that execute instructions on behalf of threads. That is, the HTP core **402** can include an instruction processing block. The HTP core **402** can further include, or can be coupled to, the thread controller **412**. The thread controller **412** can provide thread control and state for each active thread within the HTP core **402**. The data cache **406** can include cache for a host processor (e.g., for local and remote memory-compute devices, including for the HTP core **402**), and the instruction cache **404** can include cache for use by the HTP core **402**. In an example, the data cache **406** can be configured for read and write operations, and the instruction cache **404** can be configured for read only operations.

(78) In an example, the data cache **406** is a small cache provided per hardware thread. The data cache **406** can temporarily store data for use by the owning thread. The data cache **406** can be managed by hardware or software in the HTP accelerator **400**. For example, hardware can be configured to automatically allocate or evict lines as needed, as load and store operations are executed by the HTP core **402**. Software, such as using RISC-V instructions, can determine which memory accesses should be cached, and when lines should be invalidated or written back to other memory locations.

(79) Data caching on the HTP accelerator **400** has various benefits, including making larger accesses more efficient for the memory controller, allowing an executing thread to avoid stalling. However, there are situations when using the cache causes inefficiencies. An example includes accesses where data is accessed only once, and causes thrashing of the cache lines. To help address this problem, the HTP accelerator **400** can use a set of custom load instructions to force a load instruction to check for a cache hit, and on a cache miss to issue a memory request for the requested operand and not put the obtained data in the data cache **406**. The HTP accelerator **400** thus includes various different types of load instructions, including non-cached and cache line loads. The non-cached load instructions use the cached data if dirty data is present in the cache. The non-cached load instructions ignore clean data in the cache, and do not write accessed data to the data cache. For cache line load instructions, the complete data cache line (e.g., comprising 64 bytes) can be loaded from memory into the data cache **406**, and can load the addressed memory

into a specified register. These loads can use the cached data if clean or dirty data is in the data cache **406**. If the referenced memory location is not in the data cache **406**, then the entire cache line can be accessed from memory. Use of the cache line load instructions can reduce cache misses when sequential memory locations are being referenced (such as memory copy operations) but can also waste memory and bandwidth at the NOC interface **416** if the referenced memory data is not used.

(80) In an example, the HTP accelerator **400** includes a custom store instruction that is non-cached. The non-cached store instruction can help avoid thrashing the data cache **406** with write data that is not sequentially written to memory.

(81) In an example, the HTP accelerator **400** further includes a translation block **408**. The translation block **408** can include a virtual-to-physical translation block for local memory of a memory-compute device. For example, a host processor, such as in the HTP core **402**, can execute a load or store instruction, and the instruction can generate a virtual address. The virtual address can be translated to a physical address of the host processor, such as using a translation table from the translation block **408**. The memory interface **410**, for example, can include an interface between the HTP core **402** and the NOC interface **416**.

(82) FIG. 5 illustrates an example of a representation of a hybrid threading fabric (HTF), or HTF **500**, of a memory-compute device, according to an embodiment. In an example, the HTF **500** can include or comprise the HTF **142** from the example of FIG. 1. The HTF **500** is a coarse-grained, reconfigurable compute fabric that can be optimized for high-level language operand types and operators (e.g., using C/C++ or other high-level language). In an example, the HTF **500** can include configurable, n-bit wide (e.g., 512-bit wide) data paths that interconnect hardened SIMD arithmetic units.

(83) In an example, the HTF **500** comprises an HTF cluster **502** that includes multiple HTF tiles, including an example tile **504**, or Tile N. Each HTF tile can include one or more compute elements with local memory and arithmetic functions. For example, each tile can include a compute pipeline with support for integer and floating-point operations. In an example, the data path, compute elements, and other infrastructure can be implemented as hardened IP to provide maximum performance while minimizing power consumption and reconfiguration time.

(84) In the example of FIG. 5, the tiles comprising the HTF cluster **502** are linearly arranged, and each tile in the cluster can be coupled to one or multiple other tiles in the HTF cluster **502**. In the example of FIG. 5, the example tile **504**, or Tile N, is coupled to four other tiles, including to a base tile **510** (e.g., Tile N-2) via the port labeled SF IN N-2, to an adjacent tile **512** (e.g., Tile N-1) via the port labeled SF IN N-1, and to a Tile N+1 via the port labeled SF IN N+1 and to a Tile N+2 via the port labeled SF IN N+2. The example tile **504** can be coupled to the same or other tiles via respective output ports, such as those labeled SF OUT N-1, SF OUT N-2, SF OUT N+1, and SF OUT N+2. In this example, the ordered list of names for the various tiles are notional indications of the positions of the tiles. In other examples, the tiles comprising the HTF cluster **502** can be arranged in a grid or other configuration, with each the similarly coupled to one or several of its nearest neighbors in the grid. Tiles that are provided at an edge of a cluster can optionally have fewer connections to neighboring tiles. For example, Tile N-2, or the base tile **510** in the example of FIG. 5, can be coupled only to the adjacent tile **512** (Tile N-1) and to the example tile **504** (Tile N). Fewer or additional inter-tile connections can similarly be used.

(85) The HTF cluster **502** can further include memory interface modules, including a first memory interface module **506**. The memory interface modules can couple the HTF cluster **502** to a NOC, such as the first NOC **118**. In an example, the memory interface modules can allow tiles within a cluster to make requests to other locations in a memory-compute system, such as in the same or different node in the system. That is, the representation of the HTF **500** can comprise a portion of a larger fabric that can be distributed across multiple nodes, such as with one or more HTF tiles or HTF clusters at each of the nodes. Requests can be made between tiles or nodes within the context

of the larger fabric.

(86) In the example of FIG. 5, the tiles in the HTF cluster **502** are coupled using a synchronous fabric (SF). The synchronous fabric can provide communication between a particular tile and its neighboring tiles in the HTF cluster **502**, as described above. Each HTF cluster **502** can further include an asynchronous fabric (AF) that can provide communication among, e.g., the tiles in the cluster, the memory interfaces in the cluster, and a dispatch interface **508** in the cluster.

(87) In an example, the synchronous fabric can exchange messages that include data and control information. The control information can include, among other things, instruction RAM address information or a thread identifier. The control information can be used to set up a data path, and a data message field can be selected as a source for the path. Generally, the control fields can be provided or received earlier, such that they can be used to configure the data path. For example, to help minimize any delay through the synchronous domain pipeline in a tile, the control information can arrive at a tile a few clock cycles before the data field. Various registers can be provided to help coordinate dataflow timing in the pipeline.

(88) In an example, each tile in the HTF cluster **502** can include multiple memories. Each memory can have the same width as the data path (e.g., 512 bits) and can have a specified depth, such as in a range of 512 to 1024 elements. The tile memories can be used to store data that supports data path operations. The stored data can include constants loaded as part of a kernel's cluster configuration, for example, or can include variables calculated as part of the data flow. In an example, the tile memories can be written from the asynchronous fabric as a data transfer from another synchronous domain, or can include a result of a load operation such as initiated by another synchronous domain. The tile memory can be read via synchronous data path instruction execution in the synchronous domain.

(89) In an example, each tile in an HTF cluster **502** can have a dedicated instruction RAM (INST RAM). In an example of an HTF cluster **502** with sixteen tiles, and instruction RAM instances with sixty-four entries, the cluster can allow algorithms to be mapped with up to 1024 multiply-shift and/or ALU operations. The various tiles can optionally be pipelined together, such as using the synchronous fabric, to allow data flow compute with minimal memory access, thus minimizing latency and reducing power consumption. In an example, the asynchronous fabric can allow memory references to proceed in parallel with computation, thereby providing more efficient streaming kernels. In an example, the various tiles can include built-in support for loop-based constructs, and can support nested looping kernels.

(90) The synchronous fabric can allow multiple tiles to be pipelined, such as without a need for data queuing. Tiles that participate in a synchronous domain can, for example, act as a single pipelined data path. A first or base tile (e.g., Tile N-2, in the example of FIG. 5) of a synchronous domain can initiate a thread of work through the pipelined tiles. The base tile can be responsible for starting work on a predefined cadence referred to herein as a Spoke Count. For example, if the Spoke Count is 3, then the base tile can initiate work every third clock cycle.

(91) In an example, the synchronous domain comprises a set of connected tiles in the HTF cluster **502**. Execution of a thread can begin at the domain's base tile and can progress from the base tile, via the synchronous fabric, to other tiles in the same domain. The base tile can provide the instruction to be executed for the first tile. The first tile can, by default, provide the same instruction for the other connected tiles to execute. However, in some examples, the base tile, or a subsequent tile, can conditionally specify or use an alternative instruction. The alternative instruction can be chosen by having the tile's data path produce a Boolean conditional value, and then can use the Boolean value to choose between an instruction set of the current tile and the alternate instruction.

(92) The asynchronous fabric can be used to perform operations that occur asynchronously relative to a synchronous domain. Each tile in the HTF cluster **502** can include an interface to the asynchronous fabric. The inbound interface can include, for example, a FIFO buffer or queue (e.g.,

AF IN QUEUE) to provide storage for message that cannot be immediately processed. Similarly, the outbound interface of the asynchronous fabric can include a FIFO buffer or queue (e.g., AF OUT QUEUE) to provide storage for messages that cannot be immediately sent out.

(93) In an example, messages in the asynchronous fabric can be classified as data messages or control messages. Data messages can include a SIMD width data value that is written to either tile memory 0 (MEM_0) or memory 1 (MEM_1). Control messages can be configured to control thread creation, to free resources, or to issue external memory references.

(94) A tile in the HTF cluster **502** can perform various compute operations for the HTF. The compute operations can be performed by configuring the data path within the tile. In an example, a tile includes two functional blocks that perform the compute operations for the tile: a Multiply and Shift Operation block (MS OP) and an Arithmetic, Logical, and Bit Operation block (ALB OP). The two blocks can be configured to perform pipelined operations such as a Multiply and Add, or a Shift and Add, among others.

(95) In an example, each instance of a memory-compute device in a system can have a complete supported instruction set for its operator blocks (e.g., MS OP and ALB OP). In this case, binary compatibility can be realized across all devices in the system. However, in some examples, it can be helpful to maintain a base set of functionality and optional instruction set classes, such as to meet various design tradeoffs, such as die size. The approach can be similar to how the RISC-V instruction set has a base set and multiple optional instruction subsets.

(96) In an example, the example tile **504** can include a Spoke RAM. The Spoke RAM can be used to specify which input (e.g., from among the four SF tile inputs and the base tile input) is the primary input for each clock cycle. The Spoke RAM read address input can originate at a counter that counts from zero to Spoke Count minus one. In an example, different spoke counts can be used on different tiles, such as within the same HTF cluster **502**, to allow a number of slices, or unique tile instances, used by an inner loop to determine the performance of a particular application or instruction set. In an example, the Spoke RAM can specify when a synchronous input is to be written to a tile memory, for instance when multiple inputs for a particular tile instruction are used and one of the inputs arrives before the others. The early-arriving input can be written to the tile memory and can be later read when all of the inputs are available. In this example, the tile memory can be accessed as a FIFO memory, and FIFO read and write pointers can be stored in a register-based memory region or structure in the tile memory.

(97) FIG. 6A and FIG. 6B illustrate generally an example of a chiplet system that can be used to implement one or more aspects of the CNM system **102**. As similarly mentioned above, a node in the CNM system **102**, or a device within a node in the CNM system **102**, can include a chiplet-based architecture or compute-near-memory (CNM) chiplet. A packaged memory-compute device can include, for example, one, two, or four CNM chiplets. The chiplets can be interconnected using high-bandwidth, low-latency interconnects such as using a CPI interface. Generally, a chiplet system is made up of discrete modules (each a “chiplet”) that are integrated on an interposer and, in many examples, are interconnected as desired through one or more established networks to provide a system with the desired functionality. The interposer and included chiplets can be packaged together to facilitate interconnection with other components of a larger system. Each chiplet can include one or more individual integrated circuits (ICs), or “chips,” potentially in combination with discrete circuit components, and can be coupled to a respective substrate to facilitate attachment to the interposer. Most or all chiplets in a system can be individually configured for communication through established networks.

(98) The configuration of chiplets as individual modules of a system is distinct from such a system being implemented on single chips that contain distinct device blocks (e.g., intellectual property (IP) blocks) on one substrate single die), such as a system-on-a-chip (SoC), or multiple discrete packaged devices integrated on a printed circuit board (PCB). In general, chiplets provide better performance (e.g., lower power consumption, reduced latency, etc.) than discrete packaged devices,

and chiplets provide greater production benefits than single die chips. These production benefits can include higher yields or reduced development costs and time.

(99) Chiplet systems can include, for example, one or more application (or processor) chiplets and one or more support chiplets. Here, the distinction between application and support chiplets is simply a reference to the likely design scenarios for the chiplet system. Thus, for example, a synthetic vision chiplet system can include, by way of example only, an application chiplet to produce the synthetic vision output along with support chiplets, such as a memory controller chiplet, a sensor interface chiplet, or a communication chiplet. In a typical use case, the synthetic vision designer can design the application chiplet and source the support chiplets from other parties. Thus, the design expenditure (e.g., in terms of time or complexity) is reduced because by avoiding the design and production of functionality embodied in the support chiplets.

(100) Chiplets also support the tight integration of IP blocks that can otherwise be difficult, such as those manufactured using different processing technologies or using different feature sizes (or utilizing different contact technologies or spacings). Thus, multiple ICs or IC assemblies, with different physical, electrical, or communication characteristics can be assembled in a modular manner to provide an assembly with various desired functionalities. Chiplet systems can also facilitate adaptation to suit needs of different larger systems into which the chiplet system will be incorporated. In an example, ICs or other assemblies can be optimized for the power, speed, or heat generation for a specific function—as can happen with sensors—can be integrated with other devices more easily than attempting to do so on a single die. Additionally, by reducing the overall size of the die, the yield for chiplets tends to be higher than that of more complex, single die devices.

(101) FIG. 6A and FIG. 6B illustrate generally an example of a chiplet system, according to an embodiment. FIG. 6A is a representation of the chiplet system **602** mounted on a peripheral board **604**, that can be connected to a broader computer system by a peripheral component interconnect express (PCIe), for example. The chiplet, system **602** includes a package substrate **606**, an interposer **608**, and four chiplets, an application chiplet **610**, a host interface chiplet **612**, a memory controller chiplet **614**, and a memory device chiplet **616**. Other systems can include many additional chiplets to provide additional functionalities as will be apparent from the following discussion. The package of the chiplet system **602** is illustrated with a lid or cover **618**, though other packaging techniques and structures for the chiplet system can be used. FIG. 6B is a block diagram labeling the components in the chiplet system for clarity.

(102) The application chiplet **610** is illustrated as including a chiplet system NOC **620** to support a chiplet network **622** for inter-chiplet communications. In example embodiments the chiplet system NOC **620** can be included on the application chiplet **610**. In an example, the first NOC **118** from the example of FIG. 1 can be defined in response to selected support chiplets (e.g., host interface chiplet **612**, memory controller chiplet **614**, and memory device chiplet **616**) thus enabling a designer to select an appropriate number or chiplet network connections or switches for the chiplet system NOC **620**. In an example, the chiplet system NOC **620** can be located on a separate chiplet, or within the interposer **608**. In examples as discussed herein, the chiplet system NOC **620** implements a chiplet protocol interface (CPI) network.

(103) In an example, the chiplet system **602** can include or comprise a portion of the first memory-compute node **104** or the memory first memory-compute device **112**. That is, the various blocks or components of the first memory-compute device **112** can include chiplets that can be mounted on the peripheral board **604**, the package substrate **606**, and the interposer **608**. The interface components of the first memory-compute device **112** can comprise, generally, the host interface chiplet **612**, the memory and memory control-related components of the first memory-compute device **112** can comprise, generally, the memory controller chiplet **614**, the various accelerator and processor components of the first memory-compute device **112** can comprise, generally, the application chiplet **610** or instances thereof, and so on.

(104) The CPI interface, such as can be used for communication between or among chiplets in a system, is a packet-based network that supports virtual channels to enable a flexible and high-speed interaction between chiplets. CPI enables bridging from intra-chiplet networks to the chiplet network **622**. For example, the Advanced eXtensible Interface (AXI) is a widely used specification to design intra-chip communications, AXI specifications, however, cover a great variety of physical design options, such as the number of physical channels, signal timing, power, etc. Within a single chip, these options are generally selected to meet design goals, such as power consumption, speed, etc. However, to achieve the flexibility of the chiplet system, an adapter, such as CPI, is used to interface between the various AXI design options that can be implemented in the various chiplets. By enabling a physical channel to virtual channel mapping and encapsulating time-based signaling with a packetized protocol, CPI bridges intra-chiplet networks across the chiplet network **622**.

(105) CPI can use a variety of different physical layers to transmit packets. The physical layer can include simple conductive connections, or can include drivers to increase the voltage, or otherwise facilitate transmitting the signals over longer distances. An example of one such a physical layer can include the Advanced Interface Bus (AIB), which in various examples, can be implemented in the interposer **608**. AIB transmits and receives data using source synchronous data transfers with a forwarded clock. Packets are transferred across the AIB at single data rate (SDR) or dual data rate (DDR) with respect to the transmitted clock. Various channel widths are supported by AIB. The channel can be configured to have a symmetrical number of transmit (TX) and receive (RX) input/outputs (I/Os), or have a non-symmetrical number of transmitters and receivers (e.g., either all transmitters or all receivers). The channel can act as an AIB principal or subordinate depending on which chiplet provides the principal clock. AIB I/O cells support three clocking modes: asynchronous (i.e. non-clocked), SDR, and DDR. In various examples, the non-clocked mode is used for clocks and some control signals. The SDR mode can use dedicated SDR only I/O cells, or dual use SDR/DDR I/O cells.

(106) In an example, CPI packet protocols (e.g., point-to-point, or routable) can use symmetrical receive and transmit I/O cells within an AIB channel. The CPI streaming protocol allows more flexible use of the AIB I/O cells. In an example, an AIB channel for streaming mode can configure the I/O cells as all TX, all RX, or half TX and half RX. CPI packet protocols can use an AIB channel in either SDR or DDR operation modes. In an example, the AIB channel is configured in increments of 80 I/O cells (i.e. 40 TX and 40 RX) for SDR mode and 40 I/O cells for DDR mode. The CPI streaming protocol can use an AIB channel in either SDR or DDR operation modes. Here, in an example, the AIB channel is in increments of 40 I/O cells for both SDR and DDR modes. In an example, each AIB channel is assigned a unique interface identifier. The identifier is used during CPI reset and initialization to determine paired AIB channels across adjacent chiplets. In an example, the interface identifier is a 20-bit value comprising a seven-bit chiplet identifier, a seven-bit column identifier, and a six-bit link identifier. The AIB physical layer transmits the interface identifier using an AIB out-of-band shift register. The 20-bit interface identifier is transferred in both directions across an AIB interface using bits **32-51** of the shift registers.

(107) AIB defines a stacked set of AIB channels as an AIB channel column. An AIB channel column has some number of AIB channels, plus an auxiliary channel. The auxiliary channel contains signals used for AIB initialization. All AIB channels (other than the auxiliary channel) within a column are of the same configuration (e.g., all TX, all RX, or half TX and half RX, as well as having the same number of data I/O signals). In an example, AIB channels are numbered in continuous increasing order starting with the AIB channel adjacent to the AUX channel. The AIB channel adjacent to the AUX is defined to be AIB channel zero.

(108) Generally, CPI interfaces on individual chiplets can include serialization-deserialization (SERDES) hardware. SERDES interconnects work well for scenarios in which high-speed signaling with low signal count are desirable. SERDES, however, can result in additional power consumption and longer latencies for multiplexing and demultiplexing, error detection or correction

(e.g., using block level cyclic redundancy checking (CRC)), link-level retry, or forward error correction. However, when low latency or energy consumption is a primary concern for ultra-short reach, chiplet-to-chiplet interconnects, a parallel interface with clock rates that allow data transfer with minimal latency can be utilized. CPI includes elements to minimize both latency and energy consumption in these ultra-short reach chiplet interconnects.

(109) For flow control, CPI employs a credit-based technique. A recipient, such as the application chiplet **610**, provides a sender, such as the memory controller chiplet **614**, with credits that represent available buffers. In an example, a CPI recipient includes a buffer for each virtual channel for a given time-unit of transmission. Thus, if the CPI recipient supports five messages in time and a single virtual channel, the recipient has five buffers arranged in five rows (e.g., one row for each unit time). If four virtual channels are supported, then the recipient has twenty buffers arranged in five rows. Each buffer holds the payload of one CPI packet.

(110) When the sender transmits to the recipient, the sender decrements the available credits based on the transmission. Once all credits for the recipient are consumed, the sender stops sending packets to the recipient. This ensures that the recipient always has an available buffer to store the transmission.

(111) As the recipient processes received packets and frees buffers, the recipient communicates the available buffer space back to the sender. This credit return can then be used by the sender allow transmitting of additional information.

(112) The example of FIG. **6A** includes a chiplet mesh network **624** that uses a direct, chiplet-to-chiplet technique without a need for the chiplet system NOC **620**. The chiplet mesh network **624** can be implemented in CPI, or another chiplet-to-chiplet protocol. The chiplet mesh network **624** generally enables a pipeline of chiplets where one chiplet serves as the interface to the pipeline while other chiplets in the pipeline interface only with themselves.

(113) Additionally, dedicated device interfaces, such as one or more industry standard memory interfaces (such as, for example, synchronous memory interfaces, such as DDR5, DDR6), can be used to connect a device to a chiplet. Connection of a chiplet system or individual chiplets to external devices (such as a larger system can be through a desired interface (for example, a PCIe interface). Such an external interface can be implemented, in an example, through the host interface chiplet **612**, which in the depicted example, provides a PCIe interface external to chiplet system. Such dedicated chiplet interfaces **626** are generally employed when a convention or standard in the industry has converged on such an interface. The illustrated example of a Double Data Rate (DDR) interface connecting the memory controller chiplet **614** to a dynamic random access memory (DRAM) memory device chiplet **616** is just such an industry convention.

(114) Of the variety of possible support chiplets, the memory controller chiplet **614** is likely present in the chiplet system due to the near omnipresent use of storage for computer processing as well as sophisticated state-of-the-art for memory devices. Thus, using memory device chiplets **616** and memory controller chiplets **614** produced by others gives chiplet system designers access to robust products by sophisticated producers. Generally, the memory controller chiplet **614** provides a memory device-specific interface to read, write, or erase data. Often, the memory controller chiplet **614** can provide additional features, such as error detection, error correction, maintenance operations, or atomic operator execution. For some types of memory, maintenance operations tend to be specific to the memory device chiplet **616**, such as garbage collection in NAND flash or storage class memories, temperature adjustments (e.g., cross temperature management) in NAND flash memories. In an example, the maintenance operations can include logical-to-physical (L2P) mapping or management to provide a level of indirection between the physical and logical representation of data. In other types of memory, for example DRAM, some memory operations, such as refresh can be controlled by a host processor or of a memory controller at some times, and at other times controlled by the DRAM memory device, or by logic associated with one or more DRAM devices, such as an interface chip (in an example, a buffer).

(115) Atomic operators are a data manipulation that, for example, can be performed by the memory controller chiplet **614**. In other chiplet systems, the atomic operators can be performed by other chiplets. For example, an atomic operator of “increment” can be specified in a command by the application chiplet **610**, the command including a memory address and possibly an increment value. Upon receiving the command, the memory controller chiplet **614** retrieves a number from the specified memory address, increments the number by the amount specified in the command, and stores the result. Upon a successful completion, the memory controller chiplet **614** provides an indication of the command success to the application chiplet **610**. Atomic operators avoid transmitting the data across the chiplet mesh network **624**, resulting in lower latency execution of such commands.

(116) Atomic operators can be classified as built-in atomics or programmable (e.g., custom) atomics. Built-in atomics are a finite set of operations that are immutably implemented in hardware. Programmable atomics are small programs that can execute on a programmable atomic unit (PAU) (e.g., a custom atomic unit (CAU)) of the memory controller chiplet **614**.

(117) The memory device chiplet **616** can be, or include any combination of, volatile memory devices or non-volatile memories. Examples of volatile memory devices include, but are not limited to, random access memory (RAM)—such as DRAM) synchronous DRAM (SDRAM), graphics double data rate type 6 SDRAM (GDDR6 SDRAM), among others. Examples of non-volatile memory devices include, but are not limited to, negative-and-(NAND)-type flash memory, storage class memory (e.g., phase-change memory or memristor based technologies), ferroelectric RAM (FeRAM), among others. The illustrated example includes the memory device chiplet **616** as a chiplet, however, the device can reside elsewhere, such as in a different package on the peripheral board **604**. For many applications, multiple memory device chiplets can be provided. In an example, these memory device chiplets can each implement one or multiple storage technologies, and may include integrated compute hosts. In an example, a memory chiplet can include, multiple stacked memory die of different technologies, for example one or more static random access memory (SRAM) devices stacked or otherwise in communication with one or more dynamic random access memory (DRAM) devices. In an example, the memory controller chiplet **614** can serve to coordinate operations between multiple memory chiplets in the chiplet system **602**, for example, to use one or more memory chiplets in one or more levels of cache storage, and to use one or more additional memory chiplets as main memory. The chiplet system **602** can include multiple memory controller chiplet **614** instances, as can be used to provide memory control functionality for separate hosts, processors, sensors, networks, etc. A chiplet architecture, such as in the illustrated system, offers advantages in allowing adaptation to different memory storage technologies; and different memory interfaces, through updated chiplet configurations, such as without requiring redesign of the remainder of the system structure.

(118) FIG. 7 illustrates generally an example of a chiplet-based implementation for a memory-compute device, according to an embodiment. The example includes an implementation with four compute-near-memory, or CNM, chiplets, and each of the CNM chiplets can include or comprise portions of the first memory-compute device **112** or the first memory-compute node **104** from the example of FIG. 1. The various portions can themselves include or comprise respective chiplets. The chiplet-based implementation can include or use CPI-based intra-system communications, as similarly discussed above in the example chiplet system **602** from FIG. 6A and FIG. 6B.

(119) The example of FIG. 7 includes a first CNM package **700** comprising multiple chiplets. The first CNM package **700** includes a first chiplet **702**, a second chiplet **704**, a third chiplet **706**, and a fourth chiplet **708** coupled to a CNM NOC hub **710**. Each of the first through fourth chiplets can comprise instances of the same, or substantially the same, components or modules. For example, the chiplets can each include respective instances of an HTP accelerator, an HTF accelerator, and memory controllers for accessing internal or external memories.

(120) In the example of FIG. 7, the first chiplet **702** includes a first NOC hub edge **714** coupled to

the CNM NOC hub **710**. The other chiplets in the first CNM package **700** similarly include NOC hub edges or endpoints. The switches in the NOC hub edges facilitate intra-chiplet, or intra-chiplet-system, communications via the CNM NOC hub **710**.

(121) The first chiplet **702** can further include one or multiple memory controllers **716**. The memory controllers **716** can correspond to respective different NOC endpoint switches interfaced with the first NOC hub edge **714**. In an example, the memory controller **716** comprises the memory controller chiplet **614** or comprises the memory controller **130**, or comprises the memory subsystem **200**, or other memory-compute implementation. The memory controllers **716** can be coupled to respective different memory devices, for example including a first external memory module **712A** or a second external memory module **712B**. The external memory modules can include, e.g., GDDR6 memories that can be selectively accessed by the respective different chiplets in the system.

(122) The first chiplet **702** can further include a first HTP chiplet **718** and second HTP chiplet **720**, such as coupled to the first NOC hub edge **714** via respective different NOC endpoint switches. The HTP chiplets can correspond to HTP accelerators, such as the HTP **140** from the example of FIG. 1, or the HTP accelerator **400** from the example of FIG. 4. The HTP chiplets can communicate with the HTF chiplet **722**. The HTF chiplet **722** can correspond to an HTF accelerator, such as the HTF **142** from the example of FIG. 1, or the HTF **500** from the example of FIG. 5.

(123) The CNM NOC hub **710** can be coupled to NOC hub instances in other chiplets or other CNM packages by way of various interfaces and switches. For example, the CNM NOC hub **710** can be coupled to a CPI interface by way of multiple different NOC endpoints on the first CNM package **700**. Each of the multiple different NOC endpoints can be coupled, for example, to a different node outside of the first CNM package **700**. In an example, the CNM NOC hub **710** can be coupled to other peripherals, nodes, or devices using CTCPI or other, non-CPI protocols. For example, the first CNM package **700** can include a PCIe scale fabric interface (PCIe/SFI) or a CXL interface (CXL) configured to interface the first CNM package **700** with other devices. In an example, devices to which the first CNM package **700** is coupled using the various CPI, PCIe, CXL, or other fabric, can make up a common global address space.

(124) In the example of FIG. 7, the first CNM package **700** includes a host interface **724** (HIF) and a host processor (R5). The host interface **724** can correspond to, for example, the HIF **120** from the example of FIG. 1. The host processor, or R5, can correspond to the internal host processor **122** from the example of FIG. 1. The host interface **724** can include a PCI interface for coupling the first CNM package **700** to other external devices or systems. In an example, work can be initiated on the first CNM package **700**, or a tile cluster within the first CNM package **700**, by the host interface **724**. For example, the host interface **724** can be configured to command individual tile clusters, such as among the various chiplets in the first CNM package **700**, into and out of power/clock gate modes.

(125) FIG. 8 illustrates an example tiling of memory-compute devices, according to an embodiment. In FIG. 8, a tiled chiplet example **800** includes four instances of different compute-near-memory clusters of chiplets, where the clusters are coupled together. Each instance of a compute-near-memory chiplet can itself include one or more constituent chiplets (e.g., host processor chiplets, memory device chiplets, interface chiplets, and so on).

(126) The tiled chiplet example **800** includes, as one or multiple of its compute-near-memory (CNM) clusters, instances of the first CNM package **700** from the example of FIG. 7. For example, the tiled chiplet example **800** can include a first CNM cluster **802** that includes a first chiplet **810** (e.g., corresponding to the first chiplet **702**), a second chiplet **812** (e.g., corresponding to the second chiplet **704**), a third chiplet **814** (e.g., corresponding to the third chiplet **706**), and a fourth chiplet **816** (e.g., corresponding to the fourth chiplet **708**). The chiplets in the first CNM cluster **802** can be coupled to a common NOC hub, which in turn can be coupled to a NOC hub in an adjacent duster or dusters e.g., in a second CNM cluster **804** or a fourth CNM cluster **808**).

(127) In the example of FIG. 8, the tiled chiplet example **800** includes the first CNM cluster **802**, the second CNM cluster **804**, a third CNM cluster **806**, and the fourth CNM cluster **808**. The various different CNM chiplets can be configured in a common address space such that the chiplets can allocate and share resources across the different tiles. In an example, the chiplets the cluster can communicate with each other. For example, the first CNM cluster **802** can be communicatively coupled to the second CNM cluster **804** via an inter-chiplet CPI interface **818**, and the first CNM cluster **802** can be communicatively coupled to the fourth CNM cluster **808** via another or the same CPI interface. The second CNM cluster **804** can be communicatively coupled to the third CNM cluster **806** via the same or other CPI interface, and so on.

(128) In an example, one of the compute-near-memory chiplets in the tiled chiplet example **800** can include a host interface (e.g., corresponding to the host interface **724** from the example of FIG. 7) that is responsible for workload balancing across the tiled chiplet example **800**. The host interface can facilitate access to host-based command request queues and response queues, such as from outside of the tiled chiplet example **800**. The host interface can dispatch new threads of execution using hybrid threading processors and the hybrid threading fabric in one or more of the compute-near-memory chiplets in the tiled chiplet example **800**.

(129) A thread is started on the HTP accelerator when a CPI or CTCPI call packet is received and processed by the accelerator. A CPI or CTCPI call packet can originate from either the host interface (HIF) or an HTP accelerator. Calls originating from the HIF are issued by a host processor and are defined as a master thread. Calls that originate from an HTP are defined as a fiber thread. The major differences between master and fiber threads have to do with the restrictions of creating new fibers.

(130) Note that stack memory for all threads is pre-allocated when the application running on the host processor makes a system call indicating it needs to have access to HTP resources. Thereafter, when a thread is started on an HTP context the pre-allocated stack is available for use. The stack size for master threads and stack size for fibers can be different. Normally master threads have larger stack size than fibers.

(131) An HTP thread can execute a thread create (ETC) instruction to initiate a thread on an accelerator resource. Accelerator resources include HTP and HTF as well as generic accelerators. The generic accelerators allow interacting with accelerators that are not yet defined.

(132) The ETC instruction specifies the information that will be returned when the created thread completes. Before a thread can be created, space for the returned information must be reserved. When a thread eventually completes, the thread writes its return information into the reserved space and waits for the parent thread to execute a thread join instruction. The thread join instruction takes the returned information from the reserved space and transfers it to the thread's X register state for later usage by the parent thread. Once the thread is joined, then the reserved space is released.

(133) An ETC instruction can indicate that no return information is returned on completion of the created thread. This indication allows an HTP accelerator to create the thread without reserving space for return information. In this case, a counter is used to keep track of the number of outstanding threads with no return information.

(134) A parent thread executes a join instruction (EFJ, EFJA) to determine if a child has completed and to obtain returned results. All accelerator resources (HTP, HTF, and generic) use this common approach of returning thread completion status to the initiating parent thread. The returned state from a child thread may include: caller ID and 0, 1 or 2 64-bit return parameters. Threads may be created with the no return (nr) specified.

(135) The fiber join (EFJ) instruction joins threads that were created without the no return indication and ignores threads that were created with the no return indication. The EH instruction process a single completed child thread and writes the returned call ID and parameters to thread X registers for access by the parent thread.

(136) The Fiber Create (EFC) instruction initiates a thread on a hybrid threaded accelerator (HTA).

Up to four HTAs can exist within a CNM device including HTP and HTF.

(137) Fiber create instructions are illustrated in FIG. 9. Bits **0** to **6** are the opcode that specifies the EFC function and bits **27-31** are the function code that distinguish the EFC types (e.g., HTA0, HTA1, HTA0.BF, etc.). The bit fields of bits **7-11**, **15-19**, and **20-24** are reserved bits for future use.

(138) The EFC instructions specify several parameters for a fiber create operation. The Return Argument Count (rc) field may take the values (NR, 0, 1, 2). An HTP can only create a thread on an accelerator resource if the calling HTP has space to hold the return information. The return information includes both the Caller ID and up to two 64-bit return values. The fiber create instruction specifies the storage space required for the return information. The return information suffix options are NR, R0, R1 and R2. The NR suffix implies that no return information is to be stored (i.e., the join instruction will not provide any information). The R0 suffix implies that the caller ID will be available for a join instruction. The R1 and R2 suffixes imply that either one or two 64-bit return arguments in addition to the caller ID will be stored and available for a join instruction.

(139) The return information storage is allocated when the fiber create instruction is executed. If insufficient return information space is available, then the fiber create instruction is paused until space becomes available. Fiber create instructions with the NR suffix requires no return information storage and the number outstanding is limited to the size of a counter to track the number outstanding. The R0 suffix requires storage for the caller ID but no return values, and the R1 and R2 require both caller ID and return value space. The more information a set of created fibers requires results in smaller maximum outstanding fibers.

(140) A busy fail version of the EFC instruction is provided to allow an accelerator resource dispatcher to fail the EFC instruction if supported.

(141) The HTP dispatcher supports busy fail and will return failing status if an HTP context is not available on the target CNM device. The status of the fiber create instruction is written to X register A0.

(142) The HTF dispatcher does not support busy fail and will pause the EFC instruction until the dispatcher is able to find an HTF to start the created thread.

(143) Other accelerator resource dispatchers may choose to support busy fail or not.

(144) Example assembly: a. EFC.HTF b. EFC.HTP.C4.NR c. EFC.HTP.C2.R1 d. EFC.HTP.BF.C2.R1

(145) Once the EFC instruction completes, the thread continues at the instruction immediately following the EFC instruction.

(146) Note that master threads (i.e., thread created by a host) may execute any of the defined EFC instructions. HTP fiber threads (i.e., threads created by an HTP) may execute a busy fail type of EFC instruction or an EFC targeted at a non-HTP accelerator resource. A trap will result if a fiber attempts to execute a non-busy fail HTP fiber create.

(147) Threads may be created on an accelerator, such as an HTP **140** or HTF **142**. In this document, a thread may refer to a system or host thread, while a “fiber” may refer to a thread created at the HTP **140** or HTF **142**. In other words, host-dispatched work to HTP accelerators is referred to as master threads, and HTP-dispatched work to an HTP accelerator is referred to as a fiber thread. In many cases, the term “fiber thread” is interchangeable with the term “thread.”

(148) Multiple types of threads are supported by the HTP and HTF architecture. A thread may be designed as a return or “no return” type of thread. If the thread is of a return type, the fiber create instruction specifies the storage space for the return information. If the thread is a “no return” type of thread, then no return information storage is necessary, and the only limitation is the maximum number of available threads in the system.

(149) A busy fail instruction feature is provided to allow an accelerator resource dispatcher to fail the thread creation instruction. The HTP dispatcher supports busy fail and will return a failing status if an HTP context is not available on the target memory-compute system. Using this failing

status, a thread or fiber that attempts to spawn a child fiber to work on some of a workload is able to determine whether to continue processing a workload or divide the workload and assign a portion of the divided workload to a spawned fiber. This is described in more detail below.

(150) FIG. **10** is a pseudocode listing illustrating a divide and conquer process **1000**, according to an embodiment. A Main() function is illustrated that take an array WORK[] as a parameter. The size of the WORK[] array is saved in the local variable, SIZE_WORK. If the size of the WORK[] array is larger than a threshold value, then the WORK[] array is split into two parts that are roughly equal in size. A thread is allocated and the second half of the WORK[] array is passed as a parameter to the thread. The thread type is set to NO_RETURN, which allows the thread to terminate without returning a value, after the thread has completed its allocated work. The thread is then initiated with a Start() function.

(151) If the thread starts successfully, then the child thread executes the main procedure in a recursive manner, using the second half of the WORK[] array as its local WORK[] array.

(152) If the thread is not able to start, then the current thread may process some of the work locally. For example, up to the threshold amount of the WORK[] array may be processed and then another thread may be allocated and started.

(153) When the size of the WORK[] array is less than the threshold, then the current thread may process the WORK[] array without spawning any further threads (fibers).

(154) It is understood that the Process() function that acts on the WORK[] array is merely for illustration. Further, it is understood that the input to the Main() function (i.e., WORK[] array) is merely for illustration. The input may be an address range to process, a linked list, or other enumerated data structure. Processing the input to the Main() function may be performed using bitwise operations (e.g., a bit shift, logical bitwise making operation, or other bit operation). Further, although pseudocode in a high-level language is provided here, it is understood that the pseudocode may be used to program a hardware component, such as an application specific integrated circuit (ASIC) or other programmable unit.

(155) Further, it is understood that the threshold comparison in FIG. **9** is used for illustration. The threshold may be based on the number of operations that are executable in a single dock cycle, for instance, or some other metric of work able to be performed by an HTP in a given time.

(156) FIG. **11** is another code listing illustrating a divide and conquer process **1100**, according to an embodiment. In the code listing, fiber_init_parent() function takes a lower_bound, upper_bound, and pointer to the parent_segment array and loops in the range assigning each element of the array to -1. The function, init_parent_divide_and_conquer() takes the two bounds, pointer to the parent_segment, and the grainsize (i.e., the threshold size of work for an individual process to handle). In this code, an explicit call to emd_create_fiber_htp_c4_nr_bf() is made to create the thread where the second parameter is the recursive reference (as a function pointer) to this same function, init_parent_divide_and_conquer(). There is also the addition of bounds calculation to set up the child threads.

(157) FIG. **12** is a diagram illustrating a thread calling tree **1200**, according to an embodiment. Thread 0 receives work product with eight work units, illustrated as numerals 0 to 7. The threshold is equal to one work unit. As such, because there are eight work units, which is over the threshold, Thread 0 divides the work and spawns Thread 1 to handle half of the work. In the next cycle, Thread 0 checks again and finds that it has four work units ([0, 1, 2, 3]) assigned, which is over the threshold, and consequently spawns another thread, Thread 2. Meanwhile, Thread 1 also uses a similar evaluation and spawns Thread 3 with half of the work that was assigned to Thread 1 by Thread 0. This division of labor continues until the active threads have a single work unit (at or under the threshold amount).

(158) The spawned threads (fibers) may be configured as no return threads. In which case, when a thread completes processing of its work unit, it may simply terminate.

(159) FIG. **13** is a flow chart of an example of a method **1300** for thread execution control,

according to an embodiment. Operations of the method **1300** are performed by computer hardware, such as that described with respect to FIGS. **1-8** (e.g., host system **108**, host processor **122**, PAU thread control **304**, HTP core **402**, thread controller **412**, or dispatch interface **414**) or FIG. **14** (e.g., processing circuitry).

(160) At **1302**, a request to execute a first thread is received, the request includes a workload. The request may be received by interpreting code at a compiler that spawns a new thread, at a thread manager, or at a thread dispatcher, in various embodiments.

(161) At **1304**, the method **1300** measures, in a first thread, a workload to produce a metric. In an embodiment, the first thread is a master thread.

(162) At **1306**, the metric is compared with a threshold to determine that the metric is beyond the threshold.

(163) At **1308**, in response to the metric being beyond the threshold, the workload is divided into a set of sub-workloads consisting of a predefined number of equal parts from the workload. For instance, the workload may be divided approximately in half, such that multiple threads may process each half of the workload. Thus, in an embodiment, the predefined number is two.

(164) At **1310**, a second request to execute a second thread is created, where the second request including a first member of the set of sub-workloads. In an embodiment, the second thread is a fiber thread.

(165) In an embodiment, the request includes a busy-fail field that is set. In a further embodiment, the second request to execute the second thread fails, and in response to the second request to execute the second thread failing and the busy-fail field being set, the method **1300** continues by continuing processing the second member of the set of sub-workloads and creating a third request to execute the second thread, the third request including the first member of the set of sub-workloads. In an embodiment, the busy-fail field is bit in a chip-to-chip protocol interface (CTCPI) packet.

(166) In an embodiment, the second request includes a no-return field that is set. In a further embodiment, the no-return field is bit in a chip-to-chip protocol interface (CTCPI) packet. In another embodiment, the no-return field is used to signal that the second thread does not return a value to a stack position. In another embodiment, the no-return field releases the first thread from having to wait for the second thread to return.

(167) At **1312**, the first thread proceeds with processing a second member of the set of sub-workloads.

(168) In an embodiment, proceeding with processing the second member of the set of sub-workloads, includes the operations of measuring the second member to produce a second metric, comparing the metric with the threshold to determine that the second metric is beyond the threshold, dividing, in response to the second metric being beyond the threshold, the second member into a further set of sub-workloads consisting of the predefined number of equal parts from the second member, creating a third request to execute a third thread, the second request including a first member of the further set of sub-workloads, and proceeding with processing a second member of the further set of sub-workloads.

(169) In an embodiment, in response to when creating the second request to execute the second thread fails, the first thread is to process the second member of the set of sub-workloads in the first thread up to the threshold. In a further embodiment, the method **1300** includes the operations of repeating the creating the second request to execute the second thread after processing the second member of the set of sub-workloads up to the threshold.

(170) FIG. **13** illustrates a block diagram of an example machine **1300** with which, in which, or by which any one or more of the techniques (e.g., methodologies) discussed herein can be implemented. Examples, as described herein, can include, or can operate by, logic or a number of components, or mechanisms in the machine **1300**. Circuitry (e.g., processing circuitry) is a collection of circuits implemented in tangible entities of the machine **1300** that include hardware

(e.g., simple circuits, gates, logic, etc.). Circuitry membership can be flexible over time. Circuitries include members that can, alone or in combination, perform specified operations when operating. In an example, hardware of the circuitry can be immutably designed to carry out a specific operation (e.g., hardwired). In an example, the hardware of the circuitry can include variably connected physical components (e.g., execution units, transistors, simple circuits, etc.) including a machine readable medium physically modified (e.g., magnetically, electrically, moveable placement of invariant massed particles, etc.) to encode instructions of the specific operation. In connecting the physical components, the underlying electrical properties of a hardware constituent are changed, for example, from an insulator to a conductor or vice versa. The instructions enable embedded hardware (e.g., the execution units or a loading mechanism) to create members of the circuitry in hardware via the variable connections to carry out portions of the specific operation when in operation. Accordingly, in an example, the machine-readable medium elements are part of the circuitry or are communicatively coupled to the other components of the circuitry when the device is operating. In an example, any of the physical components can be used in more than one member of more than one circuitry. For example, under operation, execution units can be used in a first circuit of a first circuitry at one point in time and reused by a second circuit in the first circuitry, or by a third circuit in a second circuitry at a different time. Additional examples of these components with respect to the machine **1300**.

(171) In alternative embodiments, the machine **1300** can operate as a standalone device or can be connected (e.g., networked) to other machines. In a networked deployment, the machine **1300** can operate in the capacity of a server machine, a client machine, or both in server-client network environments. In an example, the machine **1300** can act as a peer machine in peer-to-peer (P2P) (or other distributed) network environment. The machine **1300** can be a personal computer (PC), a tablet PC, a set-top box (STB), a personal digital assistant (PDA), a mobile telephone, a web appliance, a network router, switch or bridge, or any machine capable of executing instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein, such as cloud computing, software as a service (SaaS), other computer cluster configurations.

(172) The machine **1300** (e.g., computer system) can include a hardware processor **1302** (e.g., a central processing unit (CPU), a graphics processing unit (GPU), a hardware processor core, or any combination thereof), a main memory **1304**, a static memory **1306** (e.g., memory or storage for firmware, microcode, a basic-input-output (BIOS), unified extensible firmware interface (UEFI), etc.), and mass storage device **1308** (e.g., hard drives, tape drives, flash storage, or other block devices) some or all of which can communicate with each other via an interlink **1330** (e.g., bus). The machine **1300** can further include a display device **1310**, an alphanumeric input device **1312** (e.g., a keyboard), and a user interface (UI) Navigation device **1314** (e.g., a mouse). In an example, the display device **1310**, the input device **1312**, and the UI navigation device **1314** can be a touch screen display. The machine **1300** can additionally include a mass storage device **1308** (e.g., a drive unit), a signal generation device **1318** (e.g., a speaker), a network interface device **1320**, and one or more sensor(s) **1316**, such as a global positioning system (GPS) sensor, compass, accelerometer, or other sensor. The machine **1300** can include an output controller **1328**, such as a serial (e.g., universal serial bus (USB), parallel, or other wired or wireless (e.g., infrared (IR), near field communication (NFC), etc.) connection to communicate or control one or more peripheral devices (e.g., a printer, card reader, etc.).

(173) Registers of the hardware processor **1302**, the main memory **1304**, the static memory **1306**, or the mass storage device **1308** can be, or include, a machine-readable media **1322** on which is stored one or more sets of data structures or instructions **1324** (e.g., software) embodying or used by any one or more of the techniques or functions described herein. The instructions **1324** can also

reside, completely or at least partially, within any of registers of the hardware processor **1302**, the main memory **1304**, the static memory **1306**, or the mass storage device **1308** during execution thereof by the machine **1300**. In an example, one or any combination of the hardware processor **1302**, the main memory **1304**, the static memory **1306**, or the mass storage device **1308** can constitute the machine-readable media **1322**. While the machine-readable media **1322** is illustrated as a single medium, the term “machine-readable medium” can include a single medium or multiple media (e.g., a centralized or distributed database, or associated caches and servers) configured to store the one or more instructions **1324**.

(174) The term “machine readable medium” can include any medium that is capable of storing, encoding, or carrying instructions for execution by the machine **1300** and that cause the machine **1300** to perform any one or more of the techniques of the present disclosure, or that is capable of storing, encoding or carrying data structures used by or associated with such instructions. Non-limiting machine-readable medium examples can include solid-state memories, optical media, magnetic media, and signals (e.g., radio frequency signals, other photon-based signals, sound signals, etc.). In an example, a non-transitory machine-readable medium comprises a machine-readable medium with a plurality of particles having invariant (e.g., rest) mass, and thus are compositions of matter. Accordingly, non-transitory machine-readable media are machine readable media that do not include transitory propagating signals. Specific examples of non-transitory machine readable media can include: non-volatile memory, such as semiconductor memory devices (e.g., electrically programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM)) and flash memory devices; magnetic disks, such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks.

(175) In an example, information stored or otherwise provided on the machine-readable media **1322** can be representative of the instructions **1324**, such as instructions **1324** themselves or a format from which the instructions **1324** can be derived. This format from which the instructions **1324** can be derived can include source code, encoded instructions (e.g., in compressed or encrypted form), packaged instructions (e.g., split into multiple packages), or the like. The information representative of the instructions **1324** in the machine-readable media **1322** can be processed by processing circuitry into the instructions to implement any of the operations discussed herein. For example, deriving the instructions **1324** from the information (e.g., processing by the processing circuitry) can include: compiling (e.g., from source code, object code, etc.), interpreting, loading, organizing (e.g., dynamically or statically linking), encoding, decoding, encrypting, unencrypting, packaging, unpackaging, or otherwise manipulating the information into the instructions **1324**.

(176) In an example, the derivation of the instructions **1324** can include assembly, compilation, or interpretation of the information (e.g., by the processing circuitry) to create the instructions **1324** from some intermediate or preprocessed format provided by the machine-readable media **1322**. The information, when provided in multiple parts, can be combined, unpacked, and modified to create the instructions **1324**. For example, the information can be in multiple compressed source code packages (or object code, or binary executable code, etc.) on one or several remote servers. The source code packages can be encrypted when in transit over a network and decrypted, uncompressed, assembled (e.g., linked) if necessary, and compiled or interpreted (e.g., into a library, stand-alone executable etc.) at a local machine, and executed by the local machine.

(177) The instructions **1324** can be further transmitted or received over a communications network **1326** using a transmission medium via the network interface device **1320** utilizing any one of a number of transfer protocols (e.g., frame relay, internet protocol (IP), transmission control protocol (TCP), user datagram protocol (UDP), hypertext transfer protocol (HTTP), etc.). Example communication networks can include a local area network (LAN), a wide area network (WAN), a packet data network (e.g., the Internet), mobile telephone networks (e.g., cellular networks), plain old telephone (POTS) networks, and wireless data networks (e.g., Institute of Electrical and

Electronics Engineers (IEEE) 802.11 family of standards known as Wi-Fi®, IEEE 802.16 family of standards known as WiMax®, IEEE 802.15.4 family of standards, peer-to-peer (P2P) networks, among others. In an example, the network interface device **1320** can include one or more physical jacks (e.g., Ethernet, coaxial, or phone jacks) or one or more antennas to connect to the network **1326**. In an example, the network interface device **1320** can include a plurality of antennas to wirelessly communicate using at least one of single-input multiple-output (SIMO), multiple-input multiple-output (MIMO), or multiple-input single-output (MISO) techniques. The term “transmission medium” shall be taken to include any intangible medium that is capable of storing, encoding or carrying instructions for execution by the machine **1300**, and includes digital or analog communications signals or other intangible medium to facilitate communication of such software. A transmission medium is a machine readable medium.

(178) To better illustrate the methods and apparatuses described herein, a non-limiting set of Example embodiments are set forth below as numerically identified Examples.

(179) Example 1 is an apparatus comprising: a hardware interface configured to receive a first request to execute a first thread, the first request including an indication of a workload; and processing circuitry configured to: determine the workload to produce a metric based at least in part on the indication; compare the metric with a threshold to determine that the metric is beyond the threshold; divide, based at least in part on the comparison, the workload into a set of sub-workloads consisting of predefined number of equal parts from the workload; create a second request to execute a second thread, the second request including a first member of the set of sub-workloads; and process a second member of the set of sub-workloads in the first thread.

(180) In Example 2, the subject matter of Example 1 includes, wherein the predefined number is two.

(181) In Example 3, the subject matter of Examples 1-2 includes, wherein the first thread is a master thread.

(182) In Example 4, the subject matter of Examples 1-3 includes, wherein the second thread is a fiber thread.

(183) In Example 5, the subject matter of Examples 1-4 includes, wherein the first request includes a busy-fail field that is set.

(184) In Example 6, the subject matter of Example 5 includes, wherein the second request to execute the second thread fails, and wherein, in response to the second request to execute the second thread failing and the busy-fail field being set, the processing circuitry is configured to: continue processing the second member of the set of sub-workloads; and create a third request to execute the second thread, the third request including the first member of the set of sub-workloads.

(185) In Example 7, the subject matter of Examples 5-6 includes, wherein the busy-fail field is bit in a chip-to-chip protocol interface (CTCPI) packet.

(186) In Example 8, the subject matter of Examples 1-7 includes, wherein the second request includes a no-return field that is set.

(187) In Example 9, the subject matter of Example 8 includes, wherein the no-return field is bit in a chip-to-chip protocol interface (CTCPI) packet.

(188) In Example 10, the subject matter of Examples 8-9 includes, wherein the no-return field is used to signal that the second thread does not return a value to a stack position.

(189) In Example 11, the subject matter of Examples 8-10 includes, wherein the no-return field releases the first thread from having to wait for the second thread to return.

(190) In Example 12, the subject matter of Examples 1-11 includes, wherein, to process the second member of the set of sub-workloads, the processing circuitry is configured to: determine the second member to produce a second metric; compare the second metric with the threshold to determine that the second metric is beyond the threshold; divide, based at least in part on the comparison, the second member into a further set of sub-workloads consisting of the predefined number of equal parts from the second member; create a third request to execute a third thread, the second request

including a first member of the further set of sub-workloads; and process a second member of the further set of sub-workloads.

(191) In Example 13, the subject matter of Examples 1-12 includes, wherein the operation to create the second request to execute the second thread fails, and in response, the first thread is to process the second member of the set of sub-workloads in the first thread up to the threshold.

(192) In Example 14, the subject matter of Example 13 includes, wherein the processing circuitry is to repeat the operation to create the second request to execute the second thread after processing the second member of the set of sub-workloads up to the threshold.

(193) Example 15 is a method comprising: receiving a first request to execute a first thread, the first request including an indication of a workload; determining the workload to produce a metric based at least in part on the indication; comparing the metric with a threshold to determine that the metric is beyond the threshold; dividing, based at least in part on the comparison, the workload into a set of sub-workloads consisting of predefined number of equal parts from the workload; creating a second request to execute a second thread, the second request including a first member of the set of sub-workloads; and process a second member of the set of sub-workloads in the first thread.

(194) In Example 16, the subject matter of Example 15 includes, wherein the predefined number is two.

(195) In Example 17, the subject matter of Examples 15-16 includes, wherein the first thread is a master thread.

(196) In Example 18, the subject matter of Examples 15-17 includes, wherein the second thread is a fiber thread.

(197) In Example 19, the subject matter of Examples 15-18 includes, wherein the first request includes a busy-fail field that is set.

(198) In Example 20, the subject matter of Example 19 includes, wherein the second request to execute the second thread fails, and wherein, in response to the second request to execute the second thread failing and the busy-fail field being set, the method comprises: continuing processing the second member of the set of sub-workloads; and creating a third request to execute the second thread, the third request including the first member of the set of sub-workloads.

(199) In Example 21, the subject matter of Examples 19-20 includes, wherein the busy-fail field is bit in a chip-to-chip protocol interface (CTCPI) packet.

(200) In Example 22, the subject matter of Examples 15-21 includes, wherein the second request includes a no-return field that is set.

(201) In Example 23, the subject matter of Example 22 includes, wherein the no-return field is bit in a chip-to-chip protocol interface (CTCPI) packet.

(202) In Example 24, the subject matter of Examples 22-23 includes, wherein the no-return field is used to signal that the second thread does not return a value to a stack position.

(203) In Example 25, the subject matter of Examples 22-24 includes, wherein the no-return field releases the first thread from having to wait for the second thread to return.

(204) In Example 26, the subject matter of Examples 15-25 includes, wherein, processing the second member of the set of sub-workloads, comprises: determining the second member to produce a second metric; comparing the second metric with the threshold to determine that the second metric is beyond the threshold; dividing, based at least in part on the comparison, the second member into a further set of sub-workloads consisting of the predefined number of equal parts from the second member; creating a third request to execute a third thread, the second request including a first member of the further set of sub-workloads; and processing a second member of the further set of sub-workloads.

(205) In Example 27, the subject matter of Examples 15-26 includes, wherein in response to when creating the second request to execute the second thread fails, the first thread is to process the second member of the set of sub-workloads in the first thread up to the threshold.

(206) In Example 28, the subject matter of Example 27 includes, repeating the creating the second

request to execute the second thread after processing the second member of the set of sub-workloads up to the threshold.

(207) Example 29 is an apparatus comprising: a hardware interface configured to receive a first request to execute a first thread, the first request including an indication of a workload; and processing circuitry configured to: recursively divide the workload into multiple sub-workloads and spawn child threads to process each of the multiple sub-workloads, until a sub-workload is smaller than a threshold size; and process the sub-workload that is smaller than the threshold size in a child thread.

(208) In Example 30, the subject matter of Example 29 includes, wherein workload is divided into two sub-workloads.

(209) In Example 31, the subject matter of Examples 29-30 includes, wherein the first thread is a master thread.

(210) In Example 32, the subject matter of Examples 29-31 includes, wherein the child thread is a fiber thread.

(211) In Example 33, the subject matter of Examples 29-32 includes, wherein the first request includes a busy-fail field that is set.

(212) In Example 34, the subject matter of Examples 29-33 includes, wherein the child thread that processes the sub-workload does not return a value to a calling thread.

(213) In Example 35, the subject matter of Examples 29-34 includes, wherein an instruction to execute the child thread to process the sub-workload includes a no-return field that is set.

(214) In Example 36, the subject matter of Example 35 includes, wherein the no-return field is bit in a chip-to-chip protocol interface (CTCPI) packet.

(215) In Example 37, the subject matter of Examples 35-36 includes, wherein the no-return field is used to signal that the child thread does not return a value to a stack position.

(216) In Example 38, the subject matter of Examples 35-37 includes, wherein the no-return field releases the first thread from having to wait for the child thread to return.

(217) Example 39 is at least one machine-readable medium including instructions that, when executed by processing circuitry, cause the processing circuitry to perform operations to implement any of Examples 1-38.

(218) Example 40 is an apparatus comprising means to implement any of Examples 1-38.

(219) Example 41 is a system to implement any of Examples 1-38.

(220) Example 42 is a method to implement any of Examples 1-38.

(221) The above detailed description includes references to the accompanying drawings, which form a part of the detailed description. The drawings show, by way of illustration, specific embodiments in which the invention can be practiced. These embodiments are also referred to herein as “examples”. Such examples can include elements in addition to those shown or described. However, the present inventors also contemplate examples in which only those elements shown or described are provided. Moreover, the present inventors also contemplate examples using any combination or permutation of those elements shown or described (or one or more aspects thereof), either with respect to a particular example (or one or more aspects thereof), or with respect to other examples (or one or more aspects thereof) shown or described herein.

(222) In this document, the terms “a” or “an” are used, as is common in patent documents, to include one or more than one, independent of any other instances or usages of “at least one” or “one or more.” In this document, the term “or” is used to refer to a nonexclusive or, such that “A or B” can include “A but not B,” “B but not A,” and “A and B,” unless otherwise indicated. In the appended claims, the terms “including” and “in which” are used as the plain-English equivalents of the respective terms “comprising” and “wherein”. Also, in the following claims, the terms “including” and “comprising” are open-ended, that is, a system, device, article, or process that includes elements in addition to those listed after such a term in a claim are still deemed to fall within the scope of that claim. Moreover, in the following claims, the terms “first,” “second,” and

“third,” etc. are used merely as labels, and are not intended to impose numerical requirements on their objects.

(223) The above description is intended to be illustrative, and not restrictive. For example, the above-described examples (or one or more aspects thereof) can be used in combination with each other. Other embodiments can be used, such as by one of ordinary skill in the art upon reviewing the above description. It is submitted with the understanding that it will not be used to interpret or limit the scope or meaning of the claims. Also, in the above Detailed Description, various features can be grouped together to streamline the disclosure. This should not be interpreted as intending that an unclaimed disclosed feature is essential to any claim. Rather, inventive subject matter can lie in less than all features of a particular disclosed embodiment. Thus, the following claims are hereby incorporated into the Detailed Description, with each claim standing on its own as a separate embodiment, and it is contemplated that such embodiments can be combined with each other in various combinations or permutations. The scope of the invention should be determined with reference to the appended claims, along with the full scope of equivalents to which such claims are entitled.

Claims

1. An apparatus comprising: a hardware interface configured to receive a first request to execute a first thread, the first request including an indication of a workload and a busy-fail field that is set; and processing circuitry configured to: determine the workload to produce a metric based at least in part on the indication; compare the metric with a threshold to determine that the metric is beyond the threshold; divide, based at least in part on the comparison, the workload into a set of sub-workloads consisting of a predefined number of equal parts from the workload; create a second request to execute a second thread, the second request including a first member of the set of sub-workloads; and process a second member of the set of sub-workloads in the first thread, wherein the second request to execute the second thread fails, and wherein, in response to the second request to execute the second thread failing and the busy-fail field being set, the processing circuitry is configured to: continue processing the second member of the set of sub-workloads; and create a third request to execute the second thread, the third request including the first member of the set of sub-workloads.
2. The apparatus of claim 1, wherein the predefined number is two.
3. The apparatus of claim 1, wherein the first thread is a master thread.
4. The apparatus of claim 1, wherein the second thread is a fiber thread.
5. The apparatus of claim 1, wherein the busy-fail field is bit in a chip-to-chip protocol interface (CTCPI) packet.
6. The apparatus of claim 1, wherein the second request includes a no-return field that is set.
7. The apparatus of claim 6, wherein the no-return field is bit in a chip-to-chip protocol interface (CTCPI) packet.
8. The apparatus of claim 6, wherein the no-return field is used to signal that the second thread does not return a value to a stack position.
9. The apparatus of claim 6, wherein the no-return field releases the first thread from having to wait for the second thread to return.
10. The apparatus of claim 1, wherein, to process the second member of the set of sub-workloads, the processing circuitry is configured to: determine the second member to produce a second metric; compare the second metric with the threshold to determine that the second metric is beyond the threshold; divide, based at least in part on the comparison, the second member into a further set of sub-workloads consisting of the predefined number of equal parts from the second member; create a third request to execute a third thread, the second request including a first member of the further set of sub-workloads; and process a second member of the further set of sub-workloads.

11. The apparatus of claim 1, wherein to create the second request to execute the second thread fails, and in response, the first thread is to process the second member of the set of sub-workloads in the first thread up to the threshold.
 12. The apparatus of claim 11, wherein the processing circuitry is to repeat the operation to create the second request to execute the second thread after processing the second member of the set of sub-workloads up to the threshold.
 13. A method comprising: receiving a first request to execute a first thread, the first request including an indication of a workload and a busy-fail field that is set; determining the workload to produce a metric based at least in part on the indication; comparing the metric with a threshold to determine that the metric is beyond the threshold; dividing, based at least in part on the comparison, the workload into a set of sub-workloads consisting of predefined number of equal parts from the workload; creating a second request to execute a second thread, the second request including a first member of the set of sub-workloads; and processing a second member of the set of sub-workloads in the first thread, wherein the second request to execute the second thread fails, and wherein, in response to the second request to execute the second thread failing and the busy-fail field being set, the method comprises: continuing processing the second member of the set of sub-workloads; and creating a third request to execute the second thread, the third request including the first member of the set of sub-workloads.
 14. The method of claim 13, wherein, processing the second member of the set of sub-workloads, comprises: determining the second member to produce a second metric; comparing the second metric with the threshold to determine that the second metric is beyond the threshold; dividing, based at least in part on the comparison, the second member into a further set of sub-workloads consisting of the predefined number of equal parts from the second member; creating a third request to execute a third thread, the second request including a first member of the further set of sub-workloads; and processing a second member of the further set of sub-workloads.
 15. The method of claim 13, wherein the busy-fail field is bit in a chip-to-chip protocol interface (CTCPI) packet.
 16. The method of claim 13, wherein the second request includes a no-return field that is set.
 17. The method of claim 16, wherein the no-return field is bit in a chip-to-chip protocol interface (CTCPI) packet.
 18. The method of claim 16, wherein the no-return field is used to signal that the second thread does not return a value to a stack position.
 19. The method of claim 16, wherein the no-return field releases the first thread from having to wait for the second thread to return.
-