

# US Patent & Trademark Office

## Patent Public Search | Text View

---

|                      |                        |
|----------------------|------------------------|
| United States Patent | 12386659               |
| Kind Code            | B2                     |
| Date of Patent       | August 12, 2025        |
| Inventor(s)          | Dutu; Alexandru et al. |

---

### Scheduling multiple processing-in-memory (PIM) threads and non-PIM threads

---

#### Abstract

Scheduling requests of multiple processing-in-memory threads and requests of multiple non-processing-in-memory threads is described. In accordance with the described techniques, a memory controller receives a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads from a host. The memory controller schedules an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads based on a priority associated with each of the requests and a current operating mode of the system. Requests are maintained in queues at the memory controller and are individually assigned a priority level based on time enqueued at the memory controller. Requests of a different mode than a current operating mode of the system are delayed for scheduling until at least one different mode request is escalated to a maximum priority value, at which point the memory controller initiates a system mode switch.

---

|                   |                                                                   |
|-------------------|-------------------------------------------------------------------|
| <b>Inventors:</b> | <b>Dutu; Alexandru (Kirkland, WA), Madan; Niti (Bee Cave, TX)</b> |
| <b>Applicant:</b> | <b>Advanced Micro Devices, Inc. (Santa Clara, CA)</b>             |
| <b>Family ID:</b> | <b>94126812</b>                                                   |
| <b>Assignee:</b>  | <b>Advanced Micro Devices, Inc. (Santa Clara, CA)</b>             |
| <b>Appl. No.:</b> | <b>18/214733</b>                                                  |
| <b>Filed:</b>     | <b>June 27, 2023</b>                                              |

#### Prior Publication Data

|                            |                         |
|----------------------------|-------------------------|
| <b>Document Identifier</b> | <b>Publication Date</b> |
| US 20250004826 A1          | Jan. 02, 2025           |

---

#### Publication Classification

**Int. Cl.:** G06F9/48 (20060101); G06F12/08 (20160101)

**U.S. Cl.:**

**CPC** G06F9/4881 (20130101); G06F12/08 (20130101);

**Field of Classification Search**

**USPC:** None

**References Cited**

**U.S. PATENT DOCUMENTS**

| Patent No.   | Issued Date | Patentee Name | U.S. Cl. | CPC          |
|--------------|-------------|---------------|----------|--------------|
| 12032925     | 12/2023     | Kim           | N/A      | G06F 7/523   |
| 2017/0277440 | 12/2016     | Willcock      | N/A      | G06F 12/08   |
| 2021/0110876 | 12/2020     | Seo           | N/A      | G06F 13/1668 |
| 2021/0208880 | 12/2020     | Song          | N/A      | G06F 9/30029 |
| 2024/0220107 | 12/2023     | Dutu          | N/A      | G06F 3/061   |

*Primary Examiner:* Sun; Michael

*Attorney, Agent or Firm:* FIG. 1 Patents

**Background/Summary**

**BACKGROUND**

(1) Processing-in-memory move processing of (PIM) architectures memory-intensive computations to memory. This contrasts with standard computer architectures which communicate data back and forth between a memory and a remote processing unit. In terms of data communication pathways, remote processing units of conventional computer architectures are further away from memory than processing-in-memory components. As a result, these conventional computer architectures suffer from increased data transfer latency, which can decrease overall computer performance. Further, due to the proximity to memory, PIM architectures can also provision higher memory bandwidth and reduced memory access energy relative to conventional computer architectures particularly when the volume of data transferred between the memory and the remote processing unit is large.

(2) Thus, PIM architectures enable increased computer performance while reducing data transfer latency as compared to conventional computer architectures that implement remote processing hardware. However, in PIM architectures where a memory controller receives requests from threads for both PIM operations (e.g., operations performed by a PIM component) and non-PIM operations (e.g., operations performed by a remote processing unit), extending conventional techniques for scheduling only non-PIM requests (e.g., based on static timing thresholds) to scheduling both PIM and non-PIM request threads results in degraded system performance.

**Description**

## BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram of an example system having a host with at least one core, memory hardware that includes a memory and a processing-in-memory component, and a memory controller.
- (2) FIG. 2 depicts an example of scheduling requests of multiple processing-in-memory threads with non-processing-in-memory threads.
- (3) FIG. 3 depicts an example processing-in-memory queue configured to store requests of multiple processing-in-memory threads as part of scheduling requests of multiple processing-in-memory threads and non-processing-in-memory threads.
- (4) FIG. 4 depicts an example of adjusting stall counters for non-processing-in-memory requests and processing-in-memory requests enqueued at a memory controller as part of scheduling requests of multiple processing-in-memory threads and non-processing-in-memory threads.
- (5) FIG. 5 depicts a procedure in an example implementation of scheduling requests of multiple processing-in-memory threads and non-processing-in-memory threads.
- (6) FIG. 6 depicts a procedure in an example implementation of ordering non-processing-in-memory requests and processing-in-memory requests using a multi-stage arbitration system as part of scheduling requests of multiple processing-in-memory threads and non-processing-in-memory threads.

## DETAILED DESCRIPTION

### (7) Overview

- (8) Computer architectures with PIM components implement processing devices embedded in memory hardware (e.g., memory chips). By implementing PIM components in memory hardware, PIM architectures are configured to provide memory-level processing capabilities to a variety of applications, such as applications executing on a host processing device that is communicatively coupled to the memory hardware. In such implementations where the PIM component provides memory-level processing for an application executed by the host processing device, the host processing device controls the PIM component by dispatching one or more application operations for performance to the PIM component. In some implementations, operations dispatched by the host processing device for performance by the PIM component are sent in the form of PIM requests to a memory controller. In some implementations, PIM requests are related as part of a PIM thread, where PIM requests of the PIM thread correspond to a sequence of instructions that are constrained for execution in a defined order.
- (9) In addition to scheduling requests of PIM threads for execution by the PIM component, the memory controller is tasked with scheduling non-PIM requests made by the host processing device. In contrast to PIM requests, which are executed by the PIM component using data stored in memory and thus executed by the PIM component independent of (e.g., without) traffic on a connection between memory hardware and the host processing device, executing non-PIM requests involves communicating data from memory hardware to the host processing device for processing by the host processing device. In some cases, executing non-PIM requests further involves communicating data from the host processing device back to memory hardware (e.g., transmitting a result generated by the host processing device to memory hardware for storage).
- (10) In implementations, execution of a PIM request requires bank-level parallelism in memory. Bank-level parallelism is defined as different banks of a memory channel being used to service one or more requests in parallel. For instance, in an example scenario where executing a PIM request causes a PIM component to perform an operation using data stored in a defined row of each bank in a dynamic random-access memory (DRAM) channel, bank-level parallelism is required to ensure that each bank in the DRAM channel is available before the PIM request can be executed. In contrast to PIM requests that provide bank-level parallelism in memory (e.g., PIM requests are serviced with memory commands that are sent to all banks in a DRAM channel), non-PIM requests

are serviced using memory commands issued to a single row in memory (e.g., a single row of a single DRAM channel bank).

(11) Due to the bank-level parallelism involved with executing PIM requests, scheduling PIM requests and non-PIM requests directed to a common memory channel is challenging. For instance, PIM requests cannot be executed simultaneously with executing non-PIM requests. Accordingly, a memory controller tasked with scheduling both PIM requests and non-PIM requests is also tasked with switching between two modes: one mode where PIM requests are serviced for execution by a PIM component and another mode where non-PIM requests are serviced.

(12) One challenge in transitioning from a non-PIM mode (e.g., where only non-PIM requests are executed) to a PIM mode (e.g., where only PIM requests are executed) is that all banks in a memory channel (e.g., all rows in each bank of a DRAM channel) must first be closed. This bank closure is required in order for a PIM component to service (e.g., execute) a PIM request by issuing commands on one or more rows in multiple memory channel banks. Transitioning a from a non-PIM mode to a PIM mode thus incurs a high penalty in terms of computational cost, as each bank in a memory channel needs to be pre-charged for a mode switch.

(13) For instance, in an implementation where a PIM request is executed by a PIM component of a DRAM channel, capacitors of each bank that was previously open need to be charged (e.g., by applying current to a capacitor to rewrite or otherwise restore a data value that was read from the capacitor while the bank was open). After pre-charging memory channel banks, the memory controller is able to switch to a PIM mode where a PIM component performs one or more operations (e.g., read, write, etc.) as part of executing a PIM request on one or more rows in each of the memory channel banks.

(14) Computer applications that are coded to include PIM operations typically dispatch PIM threads, and requests thereof, at irregular intervals. For instance, in an example scenario where a host processing device is running an application that includes both PIM operations and non-PIM operations, the host processing device sends a cadence of non-PIM requests to a memory controller and interleaves the non-PIM request with bursts of PIM requests for one or more PIM threads.

(15) Consequently, because there is not a constant arrival rate of PIM requests at the memory controller in this example scenario, it becomes difficult to decide when to switch between operating in a PIM mode and operating in a non-PIM mode in a manner that optimizes system performance. As described herein, optimizing system performance involves decreasing computational resource consumption, reducing system latency, reducing system power consumption, combinations thereof, and so forth.

(16) Selecting when to switch between a PIM mode and a non-PIM mode is further compounded when non-PIM requests are also received by a memory controller at irregular intervals from a host processing device. Many conventional systems implement static thresholds for switching between a PIM mode and a non-PIM mode based on a stall time of a pending request. For instance, in response to identifying that a PIM request pending at a memory controller has been stalled for a duration that approaches a defined stall time threshold for PIM requests, conventional systems switch from a non-PIM mode to a PIM mode in order to service the PIM request before its stall time satisfies (e.g., is greater than or equal to) the defined stall time threshold for PIM requests.

(17) Similarly, in response to identifying that a non-PIM request has been stalled for a duration that approaches a defined stall time threshold for non-PIM requests, conventional system memory controllers switch from a PIM mode to a non-PIM mode in order to service the non-PIM request before its stall time satisfies the defined stall time threshold for non-PIM requests. However, such conventional static stall thresholds are sub-optimal because they do not account for the irregularity of PIM requests received by a memory controller, scenarios where non-PIM requests are received by the memory controller at irregular intervals, a density of PIM requests received in a burst of PIM request by the memory controller, and so forth. Consequently, switching between PIM and non-PIM modes using conventional static thresholds results in unnecessary consumption of

computational resources, unnecessary latency in scheduling requests for execution, and overall degraded system performance.

(18) To address these problems facing conventional computer architectures with PIM components, scheduling multiple PIM threads, which each include multiple PIM requests, together with non-PIM threads is described. In implementations, a system includes memory hardware having a memory and a PIM component. The memory hardware is communicatively coupled to at least one core of at least one host, such as a core of a host processor. The system further includes a memory controller, such as a memory controller implemented locally at a host processor, a memory controller implemented at the memory module, or a memory controller implemented separate from a host processor and separate from the memory module. In implementations, the memory controller implements a scheduling module that is configured to schedule PIM requests from multiple PIM threads, along with non-PIM requests, received at the memory controller.

(19) Advantageously, the scheduling module is configured to schedule execution of PIM requests from different PIM threads in a different order than an order in which the respective different PIM threads are received at the memory controller. This represents an improvement over conventional systems that are limited to scheduling PIM requests from different PIM threads based on an order of arrival of the different PIM threads at the memory controller. To do so, the scheduling module identifies when different PIM threads access different physical addresses in the memory and efficiently schedules PIM requests from different PIM threads in a manner that optimizes system performance rather than being constrained by order of arrival at the memory controller. In doing so, the scheduling module reduces an amount of row conflicts caused by requests from different threads relative to conventional systems that schedule PIM requests based on an order of arrival at the memory controller. The scheduling module does so by assigning priority levels to individual PIM requests and individual non-PIM requests enqueued at the memory controller based on a respective time at which each request has been enqueued at the memory controller.

(20) To perform the techniques described herein, the memory controller includes a non-PIM queue and a PIM queue. The non-PIM queue is configured to enqueue non-PIM requests (e.g., DRAM commands) received from a host processor and the PIM queue is configured to enqueue PIM requests (e.g., requests to be executed by a PIM component of the memory module) received from the host processor. The scheduling module of the memory controller includes a three-stage arbitration system that schedules PIM requests and non-PIM requests for execution based on request type and request priority, with deference to a current operating mode of the system (e.g., a PIM mode where PIM requests are executed and a non-PIM mode where non-PIM requests are executed).

(21) The three-stage arbitration system is configured to consider both PIM requests and non-PIM requests at each stage of arbitration. In implementations, a first arbiter (e.g., a first arbitration stage) of the arbitration system selects different types of PIM requests from the memory controller's PIM queue and different types of non-PIM requests from the memory controller's non-PIM queue. For instance, the first arbiter selects a page hit request, a page miss request, and a page conflict request from the PIM queue. The first arbiter additionally selects a page hit request, a page miss request, and a page conflict request from the non-PIM queue. The PIM queue is organized such that PIM requests of a common thread are arranged in a dependent manner, ensuring that PIM requests for a given thread are scheduled in the proper sequence.

(22) A second arbiter (e.g., a second arbitration stage) of the arbitration system compares priorities of the requests (both PIM and non-PIM) selected by the first arbiter and outputs a request having a highest associated priority. In scenarios where requests selected by the first arbiter have the same priority, the second arbiter outputs one of the requests according to a request ranking based on request type. In some implementations, the request ranking used by the second arbiter depends on an operating mode of the system at the time the requests will be scheduled for execution (e.g., whether the system will be operating in a PIM mode or a non-PIM mode).

(23) For instance, the second arbiter is configured to identify whether a request to be scheduled will be scheduled with the system is operating in a PIM mode or a non-PIM mode. The second arbiter prioritizes scheduling requests having a same type as the current operating mode (e.g., prioritizes scheduling PIM requests when the system is operating in a PIM mode and prioritizes scheduling non-PIM requests when the system is operating in a non-PIM mode). When prioritizing requests for scheduling, the second arbiter considers priority information associated with each request, which represents a stall time of the request. In situations where, during a PIM mode, no enqueued non-PIM requests are associated with a maximum priority level, the second arbiter prioritizes selecting enqueued PIM requests. Similarly, during a non-PIM mode, when no enqueued PIM requests are associated with the maximum priority level, the second arbiter prioritizes selecting enqueued non-PIM requests.

(24) To optimize system performance, the second arbiter refrains from selecting a request of a different type than a current operating mode (e.g., at the time the selected request is scheduled to be executed by the system) until the request of the different type reaches a maximum priority level. For instance, if during a PIM mode, no enqueued PIM requests are associated with the maximum priority level and at least one enqueued non-PIM request is associated with the maximum priority level, the second arbiter initiates a mode switch and selects a maximum priority non-PIM request. Similarly, if during a non-PIM mode, no enqueued non-PIM requests are associated with the maximum priority level and at least one enqueued PIM request is at the maximum priority level, the second arbiter initiates a mode switch by scheduling the maximum priority PIM request.

(25) A third arbiter (e.g., a third arbitration stage) of the arbitration system then schedules an execution order for the requests output by the second arbiter. In implementations where the memory controller is tasked with scheduling requests for a memory channel allocated into multiple pseudo-channels (e.g., two or more pseudo-channels), the first and second arbiters are configured to perform their functionality for each pseudo-channel simultaneously. In such scenarios where the memory channel is allocated into multiple pseudo-channels, the third arbiter schedules requests output by the second arbiter (e.g., the priority winner for each of the multiple pseudo-channels) in a round-robin manner.

(26) By implementing the arbitration system and considering historical system performance metrics, the scheduling module is configured to switch the system between operating in a PIM mode and operating in a non-PIM mode based on associated priority levels and stall times during which request are enqueued at a memory controller. The scheduling policy enforced by the scheduling module is thus configured to adapt to dynamic characteristics of PIM and non-PIM operations and/or applications running concurrently in a manner that improves system performance by switching PIM and non-PIM execution modes at optimal cycles, without being constrained by an order of arrival of PIM thread operations, which is also not possible using conventional scheduling policies.

(27) Although described with respect to a single PIM component, the techniques described herein are configured for implementation by multiple processing-in-memory components in parallel (e.g., simultaneously). For instance, in an example scenario where memory is configured as DRAM, a processing-in-memory component is included at each hierarchical DRAM component (e.g., channel, bank, array, and so forth).

(28) In some aspects, the techniques described herein relate to a system including memory hardware including a memory and a processing-in-memory component, a host including at least one core, and a memory controller configured to receive, from the host, a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads, and schedule an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads based on a priority associated with each of the requests and a current operating mode of the system.

(29) In some aspects, the techniques described herein relate to a system, wherein the current

operating mode of the system is executing processing-in-memory requests and the order of execution prioritizes scheduling the requests of the plurality of processing-in-memory threads responsive to priorities associated with the requests of the plurality of non-processing-in-memory threads not satisfying a maximum priority level.

(30) In some aspects, the techniques described herein relate to a system, wherein the current operating mode of the system is executing processing-in-memory requests and the order of execution prioritizes scheduling the requests of the plurality of non-processing-in-memory threads responsive to a priority associated with at least one request of the plurality of non-processing-in-memory threads satisfying a maximum priority level.

(31) In some aspects, the techniques described herein relate to a system, wherein the current operating mode of the system is executing non-processing-in-memory requests and the order of execution prioritizes scheduling the requests of the plurality of non-processing-in-memory threads responsive to priorities associated with the requests of the plurality of processing-in-memory threads not satisfying a maximum priority level.

(32) In some aspects, the techniques described herein relate to a system, wherein the current operating mode of the system is executing non-processing-in-memory requests and the order of execution prioritizes scheduling the requests of the plurality of processing-in-memory threads responsive to a priority associated with at least one request of the plurality of processing-in-memory threads satisfying a maximum priority level.

(33) In some aspects, the techniques described herein relate to a system, wherein the memory hardware is configured to execute requests of the plurality of processing-in-memory threads by causing the processing-in-memory component to execute one or more operations using data stored in the memory.

(34) In some aspects, the techniques described herein relate to a system, wherein the host is configured to execute requests of the plurality of non-processing-in-memory threads by obtaining data stored in the memory and causing the at least one core to execute one or more operations using the data.

(35) In some aspects, the techniques described herein relate to a system, wherein the order of execution includes scheduling requests of the plurality of processing-in-memory threads in an order that is different from an order in which the requests of the plurality of processing-in-memory threads are received at the memory controller from the host.

(36) In some aspects, the techniques described herein relate to a system, wherein the memory controller is associated with a channel in the memory and the channel in the memory is allocated into two or more pseudo-channels.

(37) In some aspects, the techniques described herein relate to a system, wherein the memory controller is configured to schedule the order of execution by selecting, for each pseudo-channel of the memory, different types of requests from the plurality of processing-in-memory threads and different types of requests from the plurality of non-processing-in-memory threads using a first arbiter.

(38) In some aspects, the techniques described herein relate to a system, wherein the first arbiter selects, for each pseudo-channel of the memory, one hit request, one miss request, and one conflict request from the plurality of processing-in-memory threads.

(39) In some aspects, the techniques described herein relate to a system, wherein the first arbiter selects, for each pseudo-channel of the memory, one hit request, one miss request, and one conflict request from the plurality of non-processing-in-memory threads.

(40) In some aspects, the techniques described herein relate to a system, wherein the memory controller is configured to schedule the order of execution by selecting, using a second arbiter and for each pseudo-channel of the memory, one request from the different types of requests from the plurality of processing-in-memory threads and the different types of requests from the plurality of non-processing-in-memory threads that were selected by the first arbiter.

- (41) In some aspects, the techniques described herein relate to a system, wherein the second arbiter is configured to select the one request for each pseudo-channel of the memory based on priority information associated with each of the different types of requests from the plurality of processing-in-memory threads and the different types of requests from the plurality of non-processing-in-memory threads that were selected by the first arbiter.
- (42) In some aspects, the techniques described herein relate to a system, wherein the memory controller is configured to schedule the order of execution using a third arbiter that defines an ordering of requests selected by the second arbiter.
- (43) In some aspects, the techniques described herein relate to a method including identifying, by a memory controller of a system, that the system is operating in a processing-in-memory mode, receiving, by the memory controller from a host of the system, a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads, and scheduling, by the memory controller, an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads by prioritizing requests of the plurality of processing-in-memory threads responsive to identifying that requests of the plurality of non-processing-in-memory threads do not satisfy a maximum priority stall threshold, and prioritizing requests of the plurality of non-processing-in-memory threads responsive to identifying that at least one of the requests of the plurality of non-processing-in-memory threads satisfies the maximum priority stall threshold.
- (44) In some aspects, the techniques described herein relate to a method, wherein prioritizing the requests of the plurality of non-processing-in-memory threads is further performed responsive to identifying that the requests of the plurality of processing-in-memory threads do not satisfy the maximum priority stall threshold.
- (45) In some aspects, the techniques described herein relate to a method, further including incrementing stall counters for requests enqueued at a request queue of the memory controller and updating priority levels of the requests enqueued at the request queue during the scheduling the order of execution.
- (46) In some aspects, the techniques described herein relate to a method including identifying, by a memory controller of a system, that the system is operating in a non-processing-in-memory mode, receiving, by the memory controller from a host of the system, a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads, and scheduling, by the memory controller, an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads by prioritizing requests of the plurality of non-processing-in-memory threads responsive to identifying that requests of the plurality of processing-in-memory threads do not satisfy a maximum priority stall threshold, and prioritizing requests of the plurality of processing-in-memory threads responsive to identifying that at least one of the requests of the plurality of processing-in-memory threads satisfies the maximum priority stall threshold.
- (47) In some aspects, the techniques described herein relate to a method, wherein prioritizing the requests of the plurality of processing-in-memory threads is further performed responsive to identifying that the requests of the plurality of non-processing-in-memory threads do not satisfy the maximum priority stall threshold.
- (48) FIG. 1 is a block diagram of a system **100** that includes a host with at least one core, memory hardware with a memory and a PIM component, and a memory controller configured to grant requests by the host for the PIM component to perform transactions. In particular, the system **100** includes host **102** and memory hardware **104**, where the host **102** and the memory hardware **104** are communicatively coupled via connection/interface **106**. In one or more implementations, the host **102** includes at least one core **108**. In some implementations, the host **102** includes multiple cores **108**. For instance, in the illustrated example of FIG. 1, host **102** is depicted as including core **108(l)** and core **108 (n)**, where n represents any integer. The memory hardware **104** includes



memory **110** and processing-in-memory component **112**.

(49) In accordance with the described techniques, the host **102** and the memory hardware **104** are coupled to one another via a wired or wireless connection, which is depicted in the illustrated example of FIG. **1** as the connection/interface **106**. Example wired connections include, but are not limited to, buses (e.g., a data bus), interconnects, traces, and planes. Examples of devices in which the system **100** is implemented include, but are not limited to, supercomputers and/or computer clusters of high-performance computing (HPC) environments, servers, personal computers, laptops, desktops, game consoles, set top boxes, tablets, smartphones, mobile devices, virtual and/or augmented reality devices, wearables, medical devices, systems on chips, and other computing devices or systems.

(50) The host **102** is an electronic circuit that performs various operations on and/or using data in the memory **110**. Examples of the host **102** and/or a core **108** of the host include, but are not limited to, a central processing unit (CPU), a graphics processing unit (GPU), a field programmable gate array (FPGA), an accelerated processing unit (APU), and a digital signal processor (DSP). For example, in one or more implementations a core **108** is a processing unit that reads and executes instructions (e.g., of a program), examples of which include to add data, to move data, and to branch data.

(51) In one or more implementations, the memory hardware **104** is a circuit board (e.g., a printed circuit board), on which the memory **110** is mounted and includes the processing-in-memory component **112**. In some variations, one or more integrated circuits of the memory **110** are mounted on the circuit board of the memory hardware **104**, and the memory hardware **104** includes one or more processing-in-memory components **112**. Examples of the memory hardware **104** include, but are not limited to, a TransFlash memory module, a single in-line memory module (SIMM), and a dual in-line memory module (DIMM). In one or more implementations, the memory hardware **104** is a single integrated circuit device that incorporates the memory **110** and the processing-in-memory component **112** on a single chip. In some examples, the memory hardware **104** is composed of multiple chips that implement the memory **110** and the processing-in-memory component **112** that are vertically (“3D”) stacked together, are placed side-by-side on an interposer or substrate, or are assembled via a combination of vertical stacking or side-by-side placement.

(52) The memory **110** is a device or system that is used to store information, such as for immediate use in a device (e.g., by a core **108** of the host **102** and/or by the processing-in-memory component **112**). In one or more implementations, the memory **110** corresponds to semiconductor memory where data is stored within memory cells on one or more integrated circuits. In at least one example, the memory **110** corresponds to or includes volatile memory, examples of which include random-access memory (RAM), dynamic random-access memory (DRAM), synchronous dynamic random-access memory (SDRAM), and static random-access memory (SRAM).

(53) In some implementations, the memory **110** corresponds to or includes a cache memory of the core **108** and/or the host **102** such as a level 1 cache, a level 2 cache, a level 3 cache, and so forth. For example, the memory **110** represents high bandwidth memory (HBM) in a 3D-stacked implementation. Alternatively or additionally, the memory **110** corresponds to or includes non-volatile memory, examples of which include solid state disks (SSD), flash memory, read-only memory (ROM), programmable read-only memory (PROM), erasable programmable read-only memory (EPROM), and electronically erasable programmable read-only memory (EEPROM). The memory **110** is thus configurable in a variety of ways without departing from the spirit or scope of the described techniques.

(54) Broadly, the processing-in-memory component **112** is configured to process processing-in-memory operations involved as part of one or more transactions (e.g., operations performed as part of servicing requests of a thread received from the core **108** via the connection/interface **106**). The processing-in-memory component **112** is representative of a processor with example processing capabilities ranging from relatively simple (e.g., an adding machine) to relatively complex (e.g., a

CPU/GPU compute core). In an example, the processing-in-memory component **112** processes requests of a thread by executing associated operations using data stored in the memory **110**.

(55) Processing-in-memory contrasts with standard computer architectures which obtain data from memory, communicate the data to a remote processing unit (e.g., a core **108** of the host **102**), and process the data using the remote processing unit (e.g., using a core **108** of the host **102** rather than the processing-in-memory component **112**). In various scenarios, the data produced by the remote processing unit as a result of processing the obtained data is written back to memory, which involves communicating the produced data over the connection/interface **106** from the remote processing unit to memory. In terms of data communication pathways, the remote processing unit (e.g., a core **108** of the host **102**) is further away from the memory **110** than the processing-in-memory component **112**, both physically and topologically. As a result, conventional computer architectures suffer from increased data transfer latency, reduced data communication bandwidth, and increased data communication energy, particularly when the volume of data transferred between the memory and the remote processing unit is large, which can also decrease overall computer performance.

(56) Thus, the processing-in-memory component **112** enables increased computer performance while reducing data transfer energy as compared to standard computer architectures that implement remote processing hardware. Further, the processing-in-memory component **112** alleviates memory performance and energy bottlenecks by moving one or more memory-intensive computations closer to the memory **110**. Although the processing-in-memory component **112** is illustrated as being disposed within the memory hardware **104**, in some examples, the described benefits of memory verification using processing-in-memory are realizable through near-memory processing implementations in which the processing-in-memory component **112** is disposed in closer proximity to the memory **110** (e.g., in terms of data communication pathways) than a core **108** of the host **102**.

(57) The system **100** is further depicted as including memory controller **114**. The memory controller **114** is configured to receive requests **116** from the host **102** (e.g., from a core **108** of the host **102**). In the illustrated example, the requests **116** represent requests of one or more PIM threads, requests of one or more non-PIM threads, or combinations thereof. Although depicted in the example system **100** as being implemented separately from the host **102**, in some implementations the memory controller is implemented locally as part of the host **102**. The memory controller **114** is further configured to schedule requests **116** for a plurality of hosts, despite being depicted in the illustrated example of FIG. 1 as serving only a single host **102**. For instance, in an example implementation the memory controller **114** schedules requests for PIM threads and requests for non-PIM threads for a plurality of different hosts, where each of the plurality of different hosts include one or more cores that requests **116** to the memory controller **114** for scheduling with the memory module.

(58) In accordance with one or more implementations, the memory controller **114** is associated with a single channel of the memory **110**. For instance, the system **100** is configured to include a plurality of different memory controllers **114**, one for each of a plurality of channels of memory **110**. The techniques described herein are thus performable using a plurality of different memory controllers to schedule requests for different channels of memory **110**. In some implementations, a single channel in memory **110** is allocated into multiple pseudo-channels. In such implementations, the memory controller **114** is configured to schedule requests **116** for different pseudo-channels of a single channel in the memory **110**.

(59) As depicted in the illustrated example of FIG. 1, the memory controller **114** includes a scheduling system **118**. The scheduling system **118** is representative of a digital circuit configured to schedule requests **116** (e.g., both requests of PIM threads and requests of non-PIM threads) for execution in a manner that optimizes performance of the system **100** (e.g., limits computational resource consumption, decreases latency, and reduces power consumption of the system **100**) when

measured over execution of the requests **116**. Advantageously, using the techniques described herein, the scheduling system **118** abstains from switching between modes (e.g., switching from executing PIM requests to executing non-PIM requests, or vice versa), until requests executed in a different mode than a current mode have been stalled for a threshold duration. As described in further detail below, the threshold duration by which a request is permitted to be stalled (e.g., enqueued at the memory controller **114**) is configurable on an application-specific basis, such that the scheduling system **118** optimizes system performance in context of applications or computational tasks being performed by the system **100**.

(60) To perform the techniques described herein, the scheduling system **118** includes a request queue **120**. The request queue **120** is configured to maintain a queue of the requests **116** received at the memory controller **114** from the host **102**. In implementations, the request queue **120** is representative of two separate queues—one queue for maintaining requests of PIM threads and another queue for maintaining requests of non-PIM threads. The scheduling system **118** is configured to schedule an order of requests **116** maintained in the request queue **120** for execution (e.g., by the processing-in-memory component **112** or by the host **102**) using the techniques described herein. As depicted in the illustrated example of FIG. 1, requests selected from the request queue **120** for execution by the scheduling system **118** are represented as scheduled requests **122**. In some implementations, the scheduling system **118** selects a single request from the request queue **120** for inclusion in the scheduled requests **122** per clock cycle of the system **100**. Alternatively, the scheduling system **118** selects multiple requests from the request queue **120** for inclusion in the scheduled requests **122** per clock cycle.

(61) To determine an ordering of the scheduled requests **122**, the scheduling system **118** employs arbitration logic **124** and stall time logic **126**. The arbitration logic **124** is configured to select, for each pseudo-channel of the memory **110**, one request from the request queue **120** for output as part of the scheduled requests **122** using a three-stage arbitration process. During the three-stage arbitration process, both PIM requests and non-PIM requests are considered at each stage of arbitration. The arbitration logic **124** considers priority information associated with each request enqueued in the request queue **120** in determining which requests to select and order for output as part of the scheduled requests **122**. The stall time logic **126** represents functionality of the scheduling system **118** to keep track of how long each of the requests **116** are enqueued at the request queue **120** and to adjust a priority associated with each request maintained in the request queue **120** based on its respective time enqueued, relative to predefined priorities and their associated stall thresholds.

(62) In addition to considering request priority, the arbitration logic **124** schedules PIM requests and non-PIM requests for execution based on request type. Advantageously, in contrast to conventional systems, the scheduling system **118** enables for requests from different PIM threads to be issued as scheduled requests **122** in a different order than their order of arrival at the memory controller **114**. Functionality of the arbitration logic **124** is described in further detail below with respect to FIG. 2, an example structure of the request queue **120** is described in further detail below with respect to FIG. 3, and functionality of the stall time logic **126** is described in further detail below with respect to FIG. 4.

(63) PIM requests included in the scheduled requests **122** are transmitted by the memory controller **114** to an operation queue **128** of the processing-in-memory component **112**. The operation queue **128** is representative of a data storage structure in the processing-in-memory component **112** that maintains an ordered list of operations scheduled for sequential execution by the processing-in-memory component **112** using data stored in memory **110**. As part of executing a PIM request, the processing-in-memory component **112** generates a result **130** that includes data generated from processing data stored in the memory **110** during one or more operations involved in executing a PIM request.

(64) Instructions included in a PIM request for outputting the result **130** are configurable in a

variety of manners. For instance, in some implementations executing a PIM request causes the processing-in-memory component **112** to communicate the result **130** to a requesting source for the transaction (e.g., the host **102**). Alternatively or additionally, in some implementations instructions included in the PIM request cause the processing-in-memory component **112** to output the result **130** to a storage location in memory **110** (e.g., to update data stored in memory **110**, for subsequent access and/or retrieval by the host **102**, and so forth). Alternatively or additionally, in some implementations instructions included in the PIM request cause the processing-in-memory component **112** to store the result **130** locally (e.g., in a register of the processing-in-memory component **112**).

(65) Because the processing-in-memory component **112** executes PIM requests on behalf of the host **102**, the processing-in-memory component **112** is configured to execute PIM requests with minimal impact on the system **100** (e.g., without invalidating caches of the system **100** or causing traffic on the connection/interface **106**). For instance, the processing-in-memory component **112** performs PIM request on the memory **110** “in the background” with respect to the host **102** and the core **108**, which frees up cycles of the host **102** and/or the core **108**, reduces memory bus traffic (e.g., reduces traffic on the connection/interface **106**), and reduces power consumption relative to performing operations at the host **102** and/or the core **108**. Notably, because the processing-in-memory component **112** is closer to the memory **110** than the core **108** of the host **102** in terms of data communication pathways, evaluating data stored in memory **110** is generally completable in a shorter amount of time using the processing-in-memory component **112** than if the evaluation were performed using the core **108** of the host **102**.

(66) Non-PIM requests included in the scheduled requests **122** include one or more commands that cause the memory hardware **104** to transmit data stored in the memory **110** (e.g., at one or more storage locations identified by a non-PIM request) to the host **102** (e.g., via the interface **106**) for further processing. In some implementations, executing a non-PIM request causes the host **102** to output a result to a storage location in memory **110** (e.g., to update data stored in memory **110**, for subsequent access and/or retrieval by the host **102** or the processing-in-memory component **112**, and so forth). Alternatively or additionally, in some implementations instructions included in the non-PIM request cause the host **102** to store a result of the non-PIM request locally (e.g., in a buffer, a register, or the like of the host **102**).

(67) FIG. 2 depicts an example **200** of adaptive scheduling requests for multiple processing-in-memory threads and multiple non-processing-in-memory threads.

(68) The example **200** is depicted as including the scheduling system **118**, the request queue **120**, and the arbitration logic **124** of the memory controller **114**. In the example **200**, the scheduling system **118** is depicted as receiving at least one PIM thread **202** that includes at least one PIM request **204**. The scheduling system **118** is further depicted as receiving at least one non-PIM thread **206** that includes at least one non-PIM request **208**. In this manner, the PIM request **204** and the non-PIM request **208** are each example instances of the requests **116**.

(69) As described herein, requests of a PIM thread (e.g., PIM thread **202**) are required to be scheduled sequentially with respect to other requests of the same PIM thread. Similarly, requests of a non-PIM thread (e.g., non-PIM thread **206**) are required to be scheduled sequentially with respect to other requests of the same non-PIM thread. The techniques described herein, however, enable the scheduling system **118** to schedule requests from different PIM threads and/or different non-PIM threads in an order that differs from an order in which the requests **116** are received at the memory controller **114**, while ensuring that requests **116** of a common thread (e.g., PIM thread **202**) are scheduled sequentially for the thread.

(70) the request queue **120** includes a PIM queue **210** and a non-PIM queue **212**. The PIM queue **210** is configured to maintain PIM requests (e.g., PIM request **204**) received at the memory controller **114** and the non-PIM queue **212** is configured to maintain non-PIM requests (e.g., non-PIM request **208**) received at the memory controller **114**. In implementations, each request

enqueued at the request queue **120** (e.g., in the PIM queue **210** or the non-PIM queue **212**) includes information describing a thread identifier for the request.

(71) In implementations, the thread identifier associated with a request is generated by the memory controller **114** and represents a combination of a requestor identifier (e.g., an identifier of the host **102** from which the request was received) and a process identifier (e.g., a process, thread, or stream of requests with which the request is associated). By generating thread identifiers for different requests, the memory controller **114** ensures that enqueued requests for a same process/thread are scheduled in order relative to one another, while allowing for requests from different processes/threads to be scheduled out of order relative to one another (e.g., based on request priority and stall time thresholds).

(72) As depicted in the example **200**, the arbitration logic **124** is configured as a three-stage arbitration system that is configured to schedule an order of execution for requests maintained in the request queue **120**. Different stages of the arbitration logic **124** are represented in the illustrated example **200** by the first arbiter **214**, the second arbiter **216**, and the third arbiter **218**. The first arbiter **214**, the second arbiter **216**, and the third arbiter **218** work together to select requests from both the PIM queue **210** and the non-PIM queue **212** and output one of the selected requests as a scheduled request **220**, where the scheduled request **220** represents a request included in the scheduled requests **122**.

(73) In implementations, the first arbiter **214** is configured to select different types of PIM requests from the PIM queue **210** and select different types of non-PIM requests from the non-PIM queue **212**. For instance, the first arbiter **214** selects a PIM hit **222** request, a PIM miss **224** request, and a PIM conflict **226** request from the PIM queue **210**. The first arbiter **214** further selects a non-PIM hit **228** request, a non-PIM miss **230** request, and a non-PIM conflict **232** request from the non-PIM queue **212**. As described herein, a hit request (e.g., a PIM hit **222** or a non-PIM hit **228**) refers to a request that, at the time of execution, uses data stored in a memory location that is already open (e.g., a request for data stored in a DRAM row of the memory **110** that is open). A miss request (e.g., a PIM miss **224** or a non-PIM miss **230**) refers to a request that, at the time of execution, uses data stored in a memory location that is not open (e.g., a request for data stored in a closed DRAM row of the memory **110**). A conflict request (e.g., a PIM conflict **226** or a non-PIM conflict **232**) refers to a request that, at the time of execution, requires use of data that is being used by at least one other request.

(74) In some implementations, the first arbiter **214** identifies whether a request maintained in the request queue **120** will be a hit, a miss, or a conflict if scheduled at a certain clock cycle of the system **100** based on scheduled requests **122** previously ordered for execution by the memory controller **114**. In one or more implementations, the first arbiter **214** selects PIM requests from the PIM queue **210** and non-PIM requests from the non-PIM queue **212** based on a respective priority assigned to each request. For instance, the first arbiter **214** selects the PIM hit **222** request by identifying a plurality of PIM requests in the PIM queue **210** that would be hits if scheduled and selects from the plurality of PIM requests a request that has a highest priority. As described in further detail below with respect to FIG. 4, a request's priority is defined based on how long the request has been stalled at the memory controller **114** (e.g., enqueued at the request queue **120**). The first arbiter **214** is thus configured to ensure that a lower priority request of a certain type is not scheduled before a higher priority request of the certain type (e.g., ensure that a higher priority PIM miss **224** request is selected over a lower priority PIM miss **224** request).

(75) The second arbiter **216** is configured to select, from the PIM requests and non-PIM requests output by the first arbiter **214**, a priority winner **234**. The priority winner **234** represents one of the PIM hit **222** request, the PIM miss **224** request, the PIM conflict **226** request, the non-PIM hit **228** request, the non-PIM miss **230** request, or the non-PIM conflict **232** request. In implementations, the second arbiter **216** selects the priority winner **234** based on a current mode in which the system **100** is operating as well as respective priority information associated with the PIM hit **222** request,

the PIM miss **224** request, the PIM conflict **226** request, the non-PIM hit **228** request, the non-PIM miss **230** request, and the non-PIM conflict **232** request.

(76) As described in further detail below with respect to FIG. 4, the priority of a request is determined by comparing a stall time of the request, which indicates how long the request has been maintained in the request queue **120**, to at least one priority level stall threshold. For instance, in an example scenario where a request can be classified as low priority, medium priority, high priority, or urgent priority, different stall thresholds correspond to each priority level, such that the priority level of a request is escalated when it satisfies a stall threshold. In this manner, the arbitration logic **124** escalates priority information for a request as a stall time for a request satisfies (e.g., is greater than or equal to) a stall threshold for a priority level.

(77) Continuing the above example, in a scenario where available priority levels are Low, Medium, High, and Urgent, “Low” corresponds to a minimum priority level and “Urgent” corresponds to a maximum priority level for a request. In implementations where no request output by the first arbiter **214** are assigned a maximum priority level, the second arbiter **216** biases against mode switching and outputs a request having a highest priority that is executable in a current operating mode of the system **100**. For instance, if the system **100** is currently operating in a PIM mode and no request output by the first arbiter **214** has a maximum priority, the second arbiter **216** outputs a PIM request (e.g., the PIM hit **222** request, the PIM miss **224** request, or the PIM conflict **226** request) having the highest priority as the priority winner **234**. Alternatively, if the system **100** is currently operating in a non-PIM mode and no request output by the first arbiter **214** has a maximum priority, the second arbiter **216** outputs a non-PIM request (e.g., the non-PIM hit **228** request, the non-PIM miss **230** request, or the non-PIM conflict **232** request) having the highest priority as the priority winner **234**.

(78) In implementations where requests output by the first arbiter **214** are associated with the same priority and no requests output by the first arbiter **214** have a maximum priority, the second arbiter **216** selects a request for output as the priority winner **234** based on type (e.g., PIM hit, PIM miss, PIM conflict, non-PIM hit, non-PIM miss, or non-PIM conflict), constrained by a current operating mode of the system **100** (e.g., executing PIM requests or executing non-PIM requests).

(79) For instance, in an example implementation where the system **100** is operating in a PIM mode (e.g., executing PIM requests), the second arbiter **216** selects the priority winner **234** based on a request ranking that favors PIM hits over PIM misses over PIM conflicts. As a specific example, in a situation where the system **100** is operating in a PIM mode, no non-PIM requests output by the first arbiter **214** have a maximum priority, and the PIM hit **222** request, the PIM miss **224** request, and the PIM conflict **226** request each have a same priority level, the request ranking causes the second arbiter **216** to output the PIM hit **222** request as the priority winner **234**.

(80) In an example implementation where the system **100** is operating in a non-PIM mode (e.g., executing non-PIM requests), the second arbiter **216** selects the priority winner **234** based on a request ranking that favors non-PIM hits over non-PIM misses over non-PIM conflicts. As a specific example, in a situation where the system **100** is operating in a non-PIM mode, no PIM requests output by the first arbiter **214** have a maximum priority, and the non-PIM hit **228** request, the non-PIM miss **230** request, and the non-PIM conflict **232** request each have a same priority level, the request ranking causes the second arbiter **216** to output the non-PIM hit **228** request as the priority winner **234**.

(81) Conversely, in an implementation where at least one of the requests output by the first arbiter **214** is associated with a maximum priority (e.g., an “Urgent” priority), the second arbiter **216** triggers a mode switch if the maximum priority request is of a different mode type. As a specific example, in a scenario where the system **100** is operating in a PIM mode and no PIM requests output by the first arbiter **214** have a maximum priority, but one of the non-PIM requests has a maximum priority, the non-PIM request having the maximum priority is output as the priority winner **234**. In this manner, the second arbiter **216** refrains from causing a mode switch until a

request of a different mode (e.g., PIM or non-PIM) than a current operating mode of the system **100** reaches a maximum priority level. In situations where both PIM requests and non-PIM requests output by the first arbiter **214** have a maximum priority level, the second arbiter **216** biases against mode switching and continues to service requests of a mode in which the system **100** is currently operating.

(82) The third arbiter **218** adds the priority winner **234** to a priority winner queue **236** and selects priority winners from the priority winner queue **236** for output as the scheduled request **220**. In implementations where the memory controller **114** is associated with a single channel in memory that is not allocated into multiple pseudo-channels, the priority winner **234** is directly output by the third arbiter **218** as the scheduled request **220**. Alternatively, in implementations where the memory controller **114** is associated with a channel in memory **110** that is allocated into multiple pseudo-channels, the arbitration logic **124** is configured to perform the functionality described herein for each of the multiple pseudo-channels simultaneously.

(83) For instance, the first arbiter **214** selects a PIM hit **222** request, a PIM miss **224** request, a PIM conflict **226** request, a non-PIM hit **228** request, a non-PIM miss **230** request, and a non-PIM conflict **232** request for each of the multiple pseudo-channels. The second arbiter **216** then selects a priority winner **234** for each of the multiple pseudo-channels. In such an implementation where a memory channel is allocated into multiple pseudo-channels, the priority winners **234** for each of the pseudo-channels are individually output as a scheduled request **220** from the priority winner queue **236** by the third arbiter **218** in a round-robin manner.

(84) By implementing the arbitration logic **124**, the scheduling system **118** is configured to schedule requests from different PIM threads and different non-PIM threads in a different order than received by the memory controller **114** while optimizing system performance with respect to mode switching, which is not possible using conventional systems.

(85) FIG. 3 depicts an example **300** of a processing-in-memory queue configured to store requests of multiple processing-in-memory threads as part of scheduling requests of multiple processing-in-memory threads and memory requests.

(86) In the illustrated example, queue **302** represents an instance of the PIM queue **210**, organized in a manner that enables the arbitration logic **124** to schedule requests from different PIM threads and different non-PIM threads in a different order than received by the memory controller **114**. To do so, the queue **302** is organized such that dependent requests (e.g., requests of a common thread such as the PIM thread **202**) are placed in a same row while independent requests (e.g., requests of different PIM threads) are placed in different rows. Each of the queue **302** is configured to store requests of a thread in a linear order (e.g., according to a First In, First Out (FIFO) principle), such that requests of a thread are ensured for sequential scheduling. In some implementations, if requests from one thread span across multiple rows, the multiple rows can be marked as dependent and will be serviced in insertion order to maintain the ordering requirements of PIM requests.

(87) As depicted in FIG. 3, in some implementations the queue **302** is an array-based queue, such that a thread and its associated requests are stored in an array, where a “front” of the queue for the thread is represented by the first element of the array, and a “rear” of the queue for the thread is represented by the last element of the array. As requests for the thread are added to the queue **302**, they are added to the rear of the array assigned for the thread, and as elements are removed (e.g., scheduled as part of the scheduled requests **122**), the front of the array is shifted. Specifically, in the example **300**, the “front” of an array is represented by entries in the “Way 0” column while the “rear” of an array is represented by entries in the “Way n” column. In accordance with the described techniques, n represents any suitable integer, as indicated by the ellipses separating the “Way 3” column from the “Way n” column.

(88) Specifically, in the illustrated example of FIG. 1, the queue **302** is depicted as storing requests for four different threads—thread a, thread b, thread c, and thread d. Requests for thread a are represented as request **304**, request **306**, request **308**, request **310**, and request **312**. Requests for

thread b are represented as request **314**, request **316**, request **318**, request **320**, and request **322**. Requests for thread c are represented as request **324**, request **326**, request **328**, request **330**, and request **332**. Requests for thread d are represented as request **334**, request **336**, request **338**, request **340**, and request **342**.

(89) To ensure that PIM requests for a thread are not selected out of order with respect to other PIM requests of the same thread, the first arbiter **214** is configured to select requests only from the “Way 0” column of queue **302**. In this manner, the first arbiter **214** is permitted to select only from request **304**, request **314**, request **324** and request **334** in the queue **302** when outputting PIM requests to the second arbiter **216**. Requests of different threads maintained in the queue **302** are shaded to represent different request types identified by the arbitration logic **124** when outputting a scheduled request **220**.

(90) For instance, requests of thread a and thread d are not shaded to represent instances where request **304** and request **334** are each selectable for output as a PIM hit **222** request by the first arbiter **214**. Requests of thread b are shaded with a diagonal line background to represent an instance where request **314** is selectable for output as a PIM miss **224** request by the first arbiter **214**. Requests of thread c are shaded with a vertical line background to represent an instance where request **324** is selectable for output as a PIM conflict **226** request by the first arbiter **214**.

(91) In this manner, the illustrated example **300** depicts a situation where the first arbiter **214** is caused to output the request **314** as the PIM miss **224** request and output the request **324** as the PIM conflict **226** request to the second arbiter **216**. Similarly, the first arbiter **214** is caused to select one of the request **304** or the request **334** for output as the PIM hit **222** request to the second arbiter **216** for use in selecting a priority winner **234**. In implementations, the first arbiter **214** selects between the request **304** and the request **334** for output as the PIM hit **222** request based on priority information associated with the request **304** and priority information associated with the request **334** (e.g., selects the request having a higher associated priority level, or the request having a higher stall time counter if the associated priority levels are equal).

(92) In implementations where a request from queue **302** is selected for output as the scheduled request **220**, that request is removed from the queue **302** and other requests for the thread are shifted towards Way 0, such that a next sequential request for the thread occupies Way 0. As a specific example, consider a scenario where request **314** is selected for output as the scheduled request **220**. In this specific example, the queue **302** is updated by shifting request **316**, request **318**, request **320**, and request **322** towards Way 0, such that request **316** occupies Way 0, request **318** occupies Way 1, request **320** occupies Way 2, and request **322** occupies Way 3. Subsequent requests received for a thread are then inserted at a “rear” of the array assigned to the thread as they are received at the memory controller **114**, thus ensuring sequential scheduling of thread requests. Having considered an example organization of a PIM queue **210** that is useable by the arbitration logic **124** to schedule requests from different PIM threads and different non-PIM threads in a different order than received at the memory controller **114**, consider now example stall time logic useable to update a priority level assigned to each request.

(93) FIG. 4 depicts an example **400** of adjusting priority information associated with individual processing-in-memory requests and individual non-processing in memory requests as part of scheduling multiple processing-in-memory threads and memory requests.

(94) The illustrated example **400** is depicted as including the scheduling system **118** of the memory controller **114**, the request queue **120** the arbitration logic **124**, and the stall time logic **126**. The illustrated example **400** further illustrates how the stall time logic **126** is configured to define different priority levels and their associated stall time threshold. Specifically, stall time logic **126** is depicted as defining a first priority level **402**, which is assigned to a request when the request has been enqueued at the request queue **120** for a duration greater than or equal to the stall threshold **404**. Stall time logic **126** is further depicted as defining a second priority level **406**, which is assigned to a request when the request has been enqueued at the request queue **120** for a duration



greater than or equal to the stall threshold **408**. Finally, stall time logic **126** is depicted as defining a third priority level **410**, which is assigned to a request when the request has been enqueued at the request queue **120** for a duration greater than or equal to the stall threshold **412**.

(95) In an example scenario, the priority level **402**, the priority level **406**, and the priority level **410** represent progressively increasing priority levels. For instance, in this example scenario, requests **116** are assigned a minimum priority level (e.g., Low) upon receipt by the memory controller **114** (e.g., when first enqueued at the request queue **120**). Continuing this example scenario, after a request has been enqueued at the request queue **120** for a duration that satisfies the stall threshold **404**, the request is assigned the priority level **402** (e.g., Medium). Continuing this example scenario, if the request persists in request queue **120** for a duration that satisfies the stall threshold **408**, the priority level assigned to the request is escalated to the priority level **406** (e.g., High). Finally, in this example scenario, if the request remains enqueued in the request queue **120** for a duration that satisfies the stall threshold **412**, the request is escalated to the priority level **410** (e.g., Urgent), which represents a maximum priority that can be assigned to an enqueued request.

(96) In implementations, the stall threshold for each priority level is configurable in any suitable manner, and stall thresholds for different levels do not necessarily correlate to one another. For instance, in an example implementation, the stall threshold **404** is defined as 256 cycles, such that a request is escalated from a minimum priority to priority level **402** (e.g., Low to Medium priority) after being enqueued at the request queue **120** for 256 cycles of the system **100**. Continuing this example implementation, the stall threshold **408** is defined as 512 cycles, such that a request is escalated from priority level **402** to priority level **406** (e.g., Medium to High) after being enqueued at the request queue **120** for 512 cycles of the system **100**. Finally, in this example implementation, the stall threshold **412** is defined as 1024 cycles, such that a request is escalated from priority level **406** to the maximum priority level represented by priority level **410** (e.g., High to Urgent) after being enqueued at the request queue **120** for 1024 cycles of the system **100**. Although described herein with respect to these example priority levels, the stall time logic **126** is configurable in any suitable manner to define any number of different priority levels with respective stall thresholds defining any suitable duration.

(97) In some implementations, stall thresholds for different priority levels of the stall time logic **126** are designated by the memory controller **114** based on an injection rate at which requests **116** are received from the host **102**. Alternatively or additionally, the memory controller **114** assigns stall thresholds to different priority levels based on an amount of enqueued requests at the request queue **120**.

(98) The stall time logic **126** is configured to dynamically adjust priorities assigned to different requests enqueued at the request queue **120** based on the priority levels and their respective stall thresholds, represented in the illustrated example **400** as priority adjustment **414**. The priority adjustment **414** thus represents an instruction by the stall time logic **126** to increment a stall counter for each request maintained in the request queue **120** (e.g., increment by one for each cycle of the system **100**) and escalate a priority level for individual requests whose associated stall counter satisfies a corresponding priority level stall threshold as defined by the stall time logic **126**.

(99) For instance, the stall time logic **126** communicates the priority adjustment **414** to the request queue **120** for use in updating stall counters and priority levels associated with enqueued requests **116**. As an example, a PIM request **416** enqueued in the PIM thread **202** is associated with information defining a stall counter **418** and a priority **420** for the PIM request **416**. The stall counter **418** describes a duration at which the PIM request **416** has been stalled at the request queue **120** (e.g., a number of system **100** cycles, a temporal duration such as nanoseconds, and so forth). The priority **420** represent information describing a priority level, defined by the stall time logic **126** assigned to the PIM request **416**.

(100) In a similar manner, a non-PIM request **422** enqueued in the non-PIM queue **210** is associated with a stall counter **424** and a priority **426**. The stall counter **424** describes a duration at

which the non-PIM request **422** has been stalled at the request queue **120** (e.g., a number of system **100** cycles, a temporal duration such as nanoseconds, and so forth). The priority **426** represent information describing a priority level, defined by the stall time logic **126** assigned to the non-PIM request **422**. Thus, by assigning a priority level to each enqueued request, the scheduling system **118** enables the first arbiter **214** and the second arbiter **216** of the arbitration logic **124** to select from enqueued requests based on their associated priority levels. Similarly, the scheduling system **118** causes the memory controller **114** to refrain from mode switching until an enqueued request associated with a mode other than a current operating mode of the system **100** (e.g., PIM vs. non-PIM) has escalated to a maximum priority level.

(101) FIG. 5 depicts a procedure **500** in an example implementation of scheduling multiple processing-in-memory threads and memory requests.

(102) A plurality of processing-in-memory request threads and a plurality of non-processing-in-memory threads are received from a host at a memory controller (block **502**). The memory controller **114**, for instance, receives PIM thread **202** and at least one non-PIM thread **206** as part of requests **116** from the host **102**.

(103) An order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads are scheduled based on a type and a priority of each of the plurality of processing-in-memory thread requests and each of the plurality of non-processing-in-memory thread requests (block **504**). The arbitration logic **124** of the scheduling system **118**, for instance, schedules an order of the requests **116** for execution and outputs the scheduled order of requests as the scheduled requests **122**. In implementations, the arbitration logic **124** orders the requests **116** for output as the scheduled requests **122** based on a current operating mode of the system **100**. For instance, if the system **100** is operating in a PIM mode, the arbitration logic **124** prioritizes PIM requests for scheduling as the scheduled requests **122** over non-PIM requests based on a priority value associated with each PIM request (e.g., scheduling higher priority PIM requests over lower priority PIM requests). In a similar manner, if the system **100** is operating in a non-PIM mode, the arbitration logic **124** prioritizes non-PIM requests for scheduling as the scheduled requests **122** over PIM requests based on a priority value associated with each non-PIM request (e.g., scheduling higher priority non-PIM requests over lower priority non-PIM requests)/

(104) A priority level of at least one of the plurality of processing-in-memory thread requests or at least one of the plurality of non-processing-in-memory thread requests is modified based on a time stalled at the memory controller (block **506**). The stall time logic **126** of the scheduling system **118**, for instance, communicates a priority adjustment **414** to the request queue **120** that causes the request queue **120** to increment a stall counter for each enqueued request (e.g., increment stall counter **418** for PIM request **416** and increment stall counter **424** for non-PIM request **422**). In implementations where a stall counter satisfies a stall threshold for a priority level defined by the stall time logic **126**, the priority adjustment **414** causes the request queue **120** to update a priority associated with the request (e.g., modify the priority **420** for PIM request **416** to the priority level **402**, the priority level **406**, or the priority level **410**).

(105) Operation of a system implementing the memory controller is then switched from a current mode to a different mode in response to detecting that a stalled request associated with the different mode is assigned a maximum priority value (block **508**). For instance, if the system **100** is operating in a PIM mode, the arbitration logic **124** prioritizes PIM requests for scheduling as the scheduled requests **122** over non-PIM requests based on a priority value associated with each PIM request until a non-PIM request enqueued at the request queue **120** is escalated to a maximum priority value and no PIM requests enqueued at the request queue **120** have the maximum priority value. Upon a non-PIM request being enqueued with a maximum priority value, the arbitration logic **124** triggers a mode switch by scheduling non-PIM requests, beginning with non-PIM requests having the maximum priority value. In a similar manner, the arbitration logic **124**

continues to prioritize non-PIM requests for scheduling as the scheduled requests **122** over PIM requests based on a priority value associated with each non-PIM request until a PIM request enqueued at the request queue **120** is escalated to the maximum priority value and no non-PIM requests in the request queue **120** have the maximum priority value.

(106) Operation of the procedure **500** proceeds for additional requests **116** received at the memory controller **114**, as indicated by the arrow returning to block **502** from block **508**.

(107) FIG. **6** depicts a procedure **600** in an example implementation of scheduling multiple processing-in-memory threads and memory requests.

(108) Different types of processing-in-memory requests and different types of non-processing-in-memory requests enqueued at a memory controller are selected (block **602**). The arbitration logic **124**, for instance, uses the first arbiter **214** to select a PIM hit **222** request, a PIM miss **224** request, and a PIM conflict **226** request from the PIM queue **210**. The arbitration logic **124** further uses the first arbiter **214** to select a non-PIM hit **228** request, a non-PIM miss **230** request, and a non-PIM conflict **232** request from the non-PIM queue **212**. The arbitration logic **124** is constrained to selecting dependent PIM requests (e.g., PIM requests of a common PIM thread) in an order defined by the thread. Conversely, the arbitration logic **124** is permitted to select independent PIM requests (e.g., PIM requests from different PIM threads) without constraint as to when the independent PIM requests arrived at the memory controller **114**.

(109) A determination is made as to whether a system implementing the memory controller is currently operating in a processing-in-memory mode (block **604**). The memory controller **114**, for instance, identifies that, at the time when the different types of PIM requests and different types of non-PIM requests would be scheduled for execution, the processing-in-memory component **112** will be executing PIM requests enqueued at the operation queue **128** (e.g., a Yes determination at block **604**). Alternatively, the memory controller **114** identifies that, at the time when the different types of PIM requests and different types of non-PIM requests would be scheduled for execution, the host **102** will be executing non-PIM requests scheduled as part of the scheduled requests **122** (e.g., a No determination at block **604**).

(110) In response to a Yes determination at block **604**, a determination is made as to whether one of the selected different types of PIM requests is associated with a maximum priority level (block **606**). The second arbiter **216**, for instance, checks each PIM request output by the first arbiter **214** and identifies whether an associated priority (e.g., whether priority **420** of PIM request **416**) indicates a maximum priority level (e.g., priority level **410**). In response to identifying that one or more of the PIM requests selected at block **602** is associated with a maximum priority level (e.g., a Yes determination at block **606**), a PIM request selected at block **602** is selected for scheduling (block **608**). The second arbiter **216**, for instance, identifies one of the PIM requests output by the first arbiter **214** (e.g., the PIM hit **222** request, the PIM miss **224** request, or the PIM conflict **226** request) having the highest priority value and outputs the identified request as the priority winner **234**. In implementations where two or more of the PIM hit **222** request, the PIM miss **224** request, or the PIM conflict **226** request have the highest priority value, the second arbiter **216** favors selection of PIM hits over PIM misses, over PIM conflicts as the priority winner **234**.

(111) Alternatively, returning to block **606**, in response to identifying that no PIM requests output by the first arbiter **214** are associated with a maximum priority value (e.g., a No determination at block **606**), a determination is made as to whether one of the selected different types of non-PIM requests is associated with a maximum priority level (block **610**). The second arbiter **216**, for instance, checks each non-PIM request output by the first arbiter **214** and identifies whether an associated priority (e.g., whether priority **426** of non-PIM request **422**) indicates the maximum priority level (e.g., priority level **410**).

(112) If no non-PIM request output by the first arbiter **214** is associated with a maximum priority level (e.g., a No determination at block **610**), operation returns to block **608** and the scheduling system **118** continues to schedule highest priority PIM requests as the priority winner **234**.

Alternatively, if a non-PIM request output by the first arbiter **214** is associated with the maximum priority level (e.g., a Yes determination at block **610**), the non-PIM request with the highest priority level (e.g., the maximum priority level) is selected (block **612**).

(113) The second arbiter **216**, for instance, identifies one of the non-PIM requests output by the first arbiter **214** (e.g., the non-PIM hit **228** request, the non-PIM miss **230** request, or the non-PIM conflict **232** request) having the highest priority value and outputs the identified request as the priority winner **234**. In implementations where two or more of the non-PIM hit **228** request, the non-PIM miss **230** request, or the non-PIM conflict **232** request have the highest priority value, the second arbiter **216** favors selection of non-PIM hits over non-PIM misses, over non-PIM conflicts as the priority winner **234**.

(114) Alternatively, in response to a No determination at block **604** (e.g., when the system implementing the memory controller **114** is operating in a non-PIM mode), a determination is made as to whether one of the selected different types of non-PIM requests is associated with a maximum priority level (block **614**). The second arbiter **216**, for instance, checks each non-PIM request output by the first arbiter **214** and identifies whether an associated priority (e.g., priority **426** of non-PIM request **422**) indicates a maximum priority level (e.g., priority level **410**). In response to identifying that one or more non-PIM requests output by the first arbiter **214** have a maximum priority level (e.g., a Yes determination at block **614**), operation proceeds to block **612**, where the highest priority non-PIM request is selected.

(115) Alternatively, in response to identifying that no non-PIM requests output by the first arbiter **214** have a maximum priority level (e.g., a No determination at block **614**), a determination is made as to whether one of the selected different types of PIM requests is associated with the maximum priority level (block **616**). The second arbiter **216**, for instance, checks each PIM request output by the first arbiter **214** and identifies whether an associated priority (e.g., whether priority **420** of PIM request **416**) indicates a maximum priority level (e.g., priority level **410**). If no PIM request output by the first arbiter **214** is associated with a maximum priority level (e.g., a No determination at block **616**), operation returns to block **612**, where the highest priority non-PIM request is selected. Alternatively, if a PIM request output by the first arbiter **214** is associated with the maximum priority level, operation proceeds to block **608** where the PIM request output by the first arbiter **214** with the maximum priority level is selected.

(116) After selecting a PIM request at block **608** or selecting a non-PIM request at block **612** for output as the priority winner **234** for the memory pseudo-channel, stall counters for enqueued requests are incremented (block **618**). The request queue **120**, for instance, removes the request selected at block **608** from PIM queue **210**, or removes the request selected at block **612** from the non-PIM queue **212**, and increments a stall counter for each request still maintained in the request queue **120**. The request queue **120**, for instance, increments the stall counter **418** and the stall counter **424** based on a priority adjustment **414** instruction issued by stall time logic **126**. The priority adjustment **414** further causes any enqueued request having an associated stall counter that satisfies a priority level stall threshold to be assigned the priority level.

(117) After outputting a scheduled request **220** for each memory pseudo-channel and incrementing the stall counters of enqueued requests, operation of the procedure **600** continues by returning to block **602**.

(118) The example techniques described herein are merely illustrative and many variations are possible based on this disclosure. Although features and elements are described above in particular combinations, each feature or element is usable alone without the other features and elements or in various combinations with or without other features and elements.

(119) The various functional units illustrated in the figures and/or described herein (including, where appropriate, the host **102** having the core **108**, the memory hardware **104** having the memory **110** and the processing-in-memory component **112**, and the memory controller **114** having the scheduling system **118**) are implemented in any of a variety of different manners such as hardware

circuitry, software or firmware executing on a programmable processor, or any combination of two or more of hardware, software, and firmware. The methods provided are implemented in any of a variety of devices, such as a general-purpose computer, a processor, or a processor core. Suitable processors include, by way of example, a general purpose processor, a special purpose processor, a conventional processor, a digital signal processor (DSP), a graphics processing unit (GPU), a parallel accelerated processor, a plurality of microprocessors, one or more microprocessors in association with a DSP core, a controller, a microcontroller, Application Specific Integrated Circuits (ASICs), Field Programmable Gate Arrays (FPGAs) circuits, any other type of integrated circuit (IC), and/or a state machine.

(120) In one or more implementations, the methods and procedures provided herein are implemented in a computer program, software, or firmware incorporated in a non-transitory computer-readable storage medium for execution by a general-purpose computer or a processor. Examples of non-transitory computer-readable storage mediums include a read only memory (ROM), a random-access memory (RAM), a register, cache memory, semiconductor memory devices, magnetic media such as internal hard disks and removable disks, magneto-optical media, and optical media such as CD-ROM disks, and digital versatile disks (DVDs).

## Claims

1. A system comprising: memory hardware including a memory and a processing-in-memory component; a host including at least one core; and a memory controller configured to: receive, from the host, a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads; and schedule an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads based on a priority associated with each of the requests and a current operating mode of the system.
2. The system of claim 1, wherein the current operating mode of the system comprises executing a first type of request and the order of execution prioritizes scheduling the first type of request in response to a second type of request not having a maximum priority level.
3. The system of claim 1, wherein the current operating mode of the system comprises executing a first type of request and the order of execution prioritizes scheduling a second type of request in response to a second type of request having a maximum priority level.
4. The system of claim 3, wherein the first type of request comprises a processing-in-memory request and the second type of request comprises a non-processing-in-memory request.
5. The system of claim 3, wherein the first type of request comprises a non-processing-in-memory request and the second type of request comprises a processing-in-memory request.
6. The system of claim 1, wherein the memory hardware is configured to execute requests of the plurality of processing-in-memory threads by causing the processing-in-memory component to execute one or more operations using data stored in the memory.
7. The system of claim 1, wherein the host is configured to execute requests of the plurality of non-processing-in-memory threads by obtaining data stored in the memory and causing the at least one core to execute one or more operations using the data.
8. The system of claim 1, wherein the order of execution comprises scheduling requests of the plurality of processing-in-memory threads in an order that is different from an order in which the requests of the plurality of processing-in-memory threads are received at the memory controller from the host.
9. The system of claim 1, wherein the memory controller is associated with a channel in the memory and the channel in the memory is allocated into two or more pseudo-channels.
10. The system of claim 1, wherein the memory controller is configured to schedule the order of execution by selecting, for each pseudo-channel of the memory, different types of requests from the plurality of processing-in-memory threads and different types of requests from the plurality of non-

processing-in-memory threads using a first arbiter.

11. The system of claim 10, wherein the first arbiter selects, for each pseudo-channel of the memory, one hit request, one miss request, and one conflict request from the plurality of processing-in-memory threads.

12. The system of claim 10, wherein the first arbiter selects, for each pseudo-channel of the memory, one hit request, one miss request, and one conflict request from the plurality of non-processing-in-memory threads.

13. The system of claim 10, wherein the memory controller is configured to schedule the order of execution by selecting, using a second arbiter and for each pseudo-channel of the memory, one request from the different types of requests from the plurality of processing-in-memory threads and the different types of requests from the plurality of non-processing-in-memory threads that were selected by the first arbiter.

14. The system of claim 13, wherein the second arbiter is configured to select the one request for each pseudo-channel of the memory based on priority information associated with each of the different types of requests from the plurality of processing-in-memory threads and the different types of requests from the plurality of non-processing-in-memory threads that were selected by the first arbiter.

15. The system of claim 13, wherein the memory controller is configured to schedule the order of execution using a third arbiter that defines an ordering of requests selected by the second arbiter.

16. A method comprising: identifying, by a memory controller of a system, that the system is operating in a processing-in-memory mode; receiving, by the memory controller from a host of the system, a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads; and scheduling, by the memory controller, an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads by: prioritizing requests of the plurality of processing-in-memory threads responsive to identifying that requests of the plurality of non-processing-in-memory threads do not satisfy a maximum priority stall threshold; and prioritizing requests of the plurality of non-processing-in-memory threads responsive to identifying that at least one of the requests of the plurality of non-processing-in-memory threads satisfies the maximum priority stall threshold.

17. The method of claim 16, wherein prioritizing the requests of the plurality of non-processing-in-memory threads is further performed responsive to identifying that the requests of the plurality of processing-in-memory threads do not satisfy the maximum priority stall threshold.

18. The method of claim 16, further comprising incrementing stall counters for requests enqueued at a request queue of the memory controller and updating priority levels of the requests enqueued at the request queue during the scheduling the order of execution.

19. A device comprising: a memory controller configured to: identify that the device is operating in a non-processing-in-memory mode; receive a plurality of processing-in-memory threads and a plurality of non-processing-in-memory threads; and schedule an order of execution for requests of the plurality of processing-in-memory threads and requests of the plurality of non-processing-in-memory threads by: prioritizing requests of the plurality of non-processing-in-memory threads responsive to identifying that requests of the plurality of processing-in-memory threads do not satisfy a maximum priority stall threshold; and prioritizing requests of the plurality of processing-in-memory threads responsive to identifying that at least one of the requests of the plurality of processing-in-memory threads satisfies the maximum priority stall threshold.

20. The device of claim 19, wherein the memory controller prioritizes the requests of the plurality of processing-in-memory threads responsive to identifying that the requests of the plurality of non-processing-in-memory threads do not satisfy the maximum priority stall threshold.

---