

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386667
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Poothiyot; Ashok Pancily et al.

Dynamically selecting artificial intelligence models and hardware environments to execute tasks

Abstract

The present disclosure relates to systems, non-transitory computer-readable media, and methods for selecting machine-learning models and hardware environments for executing a task. In particular, in one or more embodiments, the disclosed systems select a designated machine-learning model for executing a task based on workload features of the task and task routing metrics for a plurality of machine-learning models. In addition, in one or more embodiments, the disclosed systems select a designated hardware environment for executing the task based on workload features for the task and task routing metrics for a plurality of hardware environments. In some embodiments, the disclosed systems select a fallback machine-learning model and a fallback hardware environment for executing the task if the designated machine-learning model or designated hardware environment are unavailable. Moreover, in one or more embodiments, the disclosed systems can pause and initiate tasks based on bandwidth availability.

Inventors: Poothiyot; Ashok Pancily (San Francisco, CA), Zafar; Ali (Fremont, CA), Penta; Anthony (Bellevue, WA), Voorhees; Stephen (Alamo, CA), Gasser; Tim (Austin, TX), Chang; Tsung-Hsiang (Bellevue, WA), Hulten; Geoff (Lynnwood, WA)

Applicant: Dropbox, Inc. (San Francisco, CA)

Family ID: 96432831

Assignee: Dropbox, Inc. (San Francisco, CA)

Appl. No.: 18/732305

Filed: June 03, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250238265 A1	Jul. 24, 2025

Related U.S. Application Data

us-provisional-application US 63623662 20240122

Publication Classification

Int. Cl.: G06F9/50 (20060101)

U.S. Cl.:

CPC G06F9/5027 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8938740	12/2014	Nishihara	709/229	G06F 9/50
9959146	12/2017	Ahuja et al.	N/A	N/A
10949252	12/2020	Krishnamurthy et al.	N/A	N/A
11302307	12/2021	Raghavendra et al.	N/A	N/A
11341420	12/2021	Loh et al.	N/A	N/A
11397529	12/2021	Chen	N/A	G06F 3/0653
11527786	12/2021	Budan	N/A	B60L 3/12
11544604	12/2022	Poothiyot	N/A	G06N 20/00
11609793	12/2022	Suh	N/A	N/A
11645111	12/2022	Wang	718/102	G06F 9/4881
11775867	12/2022	Jamei	N/A	N/A
11790242	12/2022	Agrawal et al.	N/A	N/A
11983574	12/2023	Ravindran	N/A	N/A
12033001	12/2023	Acharya et al.	N/A	N/A
2008/0066070	12/2007	Markov	N/A	N/A
2012/0198466	12/2011	Cherkasova	718/104	G06F 9/5066
2013/0212277	12/2012	Bodik	709/226	G06F 9/4887
2017/0090990	12/2016	Furman	N/A	G06F 9/5061
2018/0027060	12/2017	Metsch	N/A	G06F 9/544
2019/0102681	12/2018	Roberts	N/A	G06N 20/00
2019/0102700	12/2018	Babu	N/A	G06N 20/00
2019/0146848	12/2018	Rastogi	718/104	G06F 9/5061

2019/0188605	12/2018	Zavesky	N/A	G06N 20/00
2019/0197011	12/2018	Zavesky et al.	N/A	N/A
2019/0266015	12/2018	Chandra et al.	N/A	N/A
2019/0303207	12/2018	Vadapandeshwara	N/A	G06F 9/5038
2020/0012962	12/2019	Dent	N/A	G06N 20/00
2020/0372307	12/2019	Arun	N/A	G06F 18/2178
2021/0089961	12/2020	Zeise	N/A	G06N 20/00
2021/0089969	12/2020	Sengupta et al.	N/A	N/A
2021/0224585	12/2020	Schmidt	N/A	G06F 9/3836
2021/0256422	12/2020	Unterthiner et al.	N/A	N/A
2021/0279576	12/2020	Shazeer et al.	N/A	N/A
2021/0334700	12/2020	Nagaraja	N/A	N/A
2022/0035679	12/2021	Sunwoo	N/A	G06N 20/00
2022/0058477	12/2021	Hu et al.	N/A	N/A
2022/0067752	12/2021	Fang	N/A	H04L 63/20
2022/0083389	12/2021	Poothia	N/A	G06F 9/5027
2022/0114026	12/2021	Cropper	N/A	G06F 9/5027
2022/0207349	12/2021	Fusco et al.	N/A	N/A
2022/0215008	12/2021	Adibowo	N/A	G06F 12/0802
2022/0237208	12/2021	Srinivasan	N/A	G06F 16/26
2022/0318689	12/2021	Li-Bland	N/A	G06N 20/00
2022/0374784	12/2021	Sharma	N/A	G06F 18/217
2023/0139459	12/2022	Latkar	718/104	G06F 9/5011
2023/0206134	12/2022	Kennel et al.	N/A	N/A
2023/0252340	12/2022	Ramanathan	706/11	G06N 20/00
2023/0259747	12/2022	Lee et al.	N/A	N/A
2023/0333901	12/2022	Kale	N/A	G06F 9/5055
2023/0376725	12/2022	Mesmakhosroshahi et al.	N/A	N/A
2023/0385129	12/2022	Johnson et al.	N/A	N/A
2023/0409874	12/2022	Silva Tavares et al.	N/A	N/A
2024/0104407	12/2023	McCarthy	N/A	G06F 9/5027
2024/0107347	12/2023	Pirmagomedov et al.	N/A	N/A
2024/0135247	12/2023	Johnsson	N/A	G06N 20/00
2024/0144051	12/2023	Kirshenboim et al.	N/A	N/A
2024/0160491	12/2023	Taneja	N/A	G06Q 10/0631
2024/0289645	12/2023	Makhijani	N/A	G06N 5/01
2024/0406166	12/2023	Bell	N/A	G06N 3/045

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
116074179	12/2022	CN	N/A
117217201	12/2022	CN	N/A
3822865	12/2020	EP	N/A
7195365	12/2021	JP	N/A
2019104063	12/2019	WO	N/A
2023146883	12/2022	WO	N/A

OTHER PUBLICATIONS

Sharifian et al.; “A predictive and probabilistic load-balancing algorithm for cluster-based web servers”; 2010 Elsevier B.V.; doi: 10.1016/j.asoc.2010.01.017 (NPL: Sharifian_2011.pdf; pp. 970-981) (Year: 2011). cited by examiner

Gai et al.; “Optimal resource allocation using reinforcement learning for IoT content-centric services”; 2018 Elsevier B.V.; <https://doi.org/10.1016/j.asoc.2018.03.056> (Gai_2018.pdf, pp. 12-21) (Year: 2018). cited by examiner

Non-Final Office Action from U.S. Appl. No. 18/732,297, mailed Aug. 26, 2024, 20 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 18/732,320, mailed Aug. 28, 2024, 14 pages. cited by applicant

Weiss M., et al., “Fail-Safe Execution of Deep Learning based Systems through Uncertainty Monitoring,” IEEE, Feb. 1, 2021, 12 pages. cited by applicant

Final Office Action from U.S. Appl. No. 18/732,297 mailed Nov. 7, 2024, 21 pages. cited by applicant

Notice of Allowance for U.S. Appl. No. 18/732,320 mailed on Nov. 14, 2024, 7 pages. cited by applicant

Invitation to Pay Additional Fees for International Application No. PCT/US2024/059866 mailed Mar. 31, 2025, 16 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 18/732,297, mailed May 6, 2025, 37 pages. cited by applicant

Notice of Allowance for U.S. Appl. No. 18/732,320 mailed on Apr. 16, 2025, 7 pages. cited by applicant

International Search Report and Written Opinion for Application No. PCT/US2024/059866, mailed on May 22, 2025, 19 pages. cited by applicant

Primary Examiner: Patel; Hiren P

Attorney, Agent or Firm: Keller Preece PLLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application claims priority to and the benefit of U.S. Provisional Patent Application No. 63/623,662, filed on Jan. 22, 2024, which is incorporated herein by reference in its entirety.

BACKGROUND

- (1) Recent years have seen significant improvements in the capabilities and capacities of artificial intelligence models. For example, artificial intelligence systems can quickly generate generative output in response to an input. To illustrate, large language models are particularly adept at interpreting natural language prompts, adapting to context based on the prompt, and can generate a variety of different outputs, including, among other things, creating content based on existing content items, providing summaries of documents, generating code, and retrieving information. However, there are a number of technical deficiencies with regard to utilizing and optimizing among various artificial intelligence models that often require large amounts of computational resources, utilize different hardware components, and provide differing qualities of output.
- (2) For example, conventional systems are inflexible regarding artificial intelligence models. Often, conventional systems can only utilize one (or a small handful) of artificial intelligence models for executing tasks. Not only does this limit conventional systems to that model, but when developers create systems, applications, and/or integrations to interface with existing systems, these inflexibilities are amplified as developers are often limited to creating systems that utilize certain artificial intelligence models. In addition, as new and more powerful artificial intelligence models are generated, existing systems must be completely rebuilt or reworked in order to interface with each new artificial intelligence model.
- (3) Also, partly due to their inflexibility, conventional systems are computationally inefficient when selecting artificial intelligence models. For instance, as mentioned, conventional systems are often built to access a specific artificial intelligence model, and thus, conventional systems often simply execute tasks by sending tasks to that model. However, as artificial intelligence models have varying capabilities and performance characteristics, simply using a single artificial intelligence model, or even just a small handful of models, leads to computational inefficiencies. In some cases, providing an artificial intelligence model with a complex task or prompt will require excessive bandwidth usage to transmit data to and from the artificial intelligence model or make multiple API calls to break the task into smaller segments. In other cases, when conventional systems utilize more than one artificial intelligence model, conventional systems will simply provide tasks to available artificial intelligence models, resulting in large amounts of bandwidth usage and processing power to run tasks on poorly fitting artificial intelligence models. For example, conventional systems may run a basic or low-complexity task on a large artificial intelligence model, utilizing more processing power than needed to complete the low-complexity task.
- (4) In addition to their inefficiencies, conventional systems are also inflexible with the hardware environments used to execute tasks that utilize artificial intelligence models. For instance, conventional systems often simply utilize a local hardware environment to execute tasks on artificial intelligence machines and only utilize other hardware environments when the local environment exhausts resources or goes offline due to issues or outages. Moreover, when conventional systems utilize other hardware environments, they simply execute tasks in hardware environments with immediate availability, regardless of whether executing the task in that hardware environment results in a computational cost or if the quality of output for the task is compromised.
- (5) Moreover, conventional systems are inefficient in their allocation of hardware resources, particularly when allocating hardware for executing training tasks. For example, conventional systems often initiate a training task on a local hardware environment, and unexpected increases in bandwidth usage combined with high bandwidth usage for the training task result in little to no bandwidth usage for other tasks. Moreover, conventional systems often train without regard to where training data is stored, often resulting in additional computing resources and processing time to move the training data to the hardware environment in order to execute the training task. Additional features and advantages of one or more embodiments of the present disclosure are

outlined in the description that follows and, in part, will be obvious from the description or may be learned by the practice of such example embodiments.

BRIEF SUMMARY

(6) Embodiments of the present disclosure provide benefits and/or solve one or more of the foregoing or other problems in the art with systems, non-transitory computer-readable media, and methods for allocating workloads to specific artificial intelligence models and hardware infrastructure. For example, in one or more embodiments, the disclosed systems dynamically and intelligently select a machine-learning model for executing a task based on the workload features of the task and task routing metrics of a plurality of machine-learning models. In one or more embodiments, the disclosed systems select additional machine-learning models as fallback models for executing the task if the selected machine-learning model is unavailable. In addition to selecting one or more designated machine-learning models, in one or more embodiments, the disclosed systems select a hardware environment for executing the task based on the workload features of the task and task routing metrics for the hardware environment. In one or more embodiments, the disclosed systems also select a fallback hardware environment for executing the task if the selected hardware environment is unavailable.

(7) In addition, in one or more embodiments, the disclosed systems intelligently schedule and initiate training tasks based on bandwidth availability. For example, in one or more embodiments, the disclosed systems can monitor hardware usage and intelligently schedule and/or initiate training tasks. Further, in one or more embodiments, the disclosed systems initiate a training task based on bandwidth availability and pauses the training task upon detecting a change in bandwidth availability. Additional features and advantages of one or more embodiments of the present disclosure are outlined in the description that follows and, in part, will be obvious from the description or may be learned by the practice of such example embodiments.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The detailed description provides one or more embodiments with additional specificity and detail through the use of the accompanying drawings, as briefly described below.

(2) FIG. 1 illustrates a diagram of an environment in which an intelligent selection and execution platform can operate in accordance with one or more embodiments.

(3) FIG. 2 illustrates an example diagram of an overview of an intelligent selection and execution platform utilizing a smart pocket ML model in accordance with one or more embodiments.

(4) FIG. 3 illustrates an example diagram of an overview of an intelligent selection and execution platform utilizing a smart pocket machine-learning model in accordance with one or more embodiments.

(5) FIG. 4 illustrates an example diagram of an overview of layers of an intelligent selection and execution platform in accordance with one or more embodiments.

(6) FIG. 5 illustrates a schematic diagram of an overview of an intelligent selection and execution platform selecting a designated machine-learning model for executing a task in accordance with one or more embodiments.

(7) FIG. 6 illustrates a schematic diagram of an intelligent selection and execution platform extracting workload features from workload data requesting execution of a task in accordance with one or more embodiments.

(8) FIG. 7 illustrates a schematic diagram of an intelligent selection and execution platform determining task routing metrics for a plurality of machine-learning models hosted in respective network environments in accordance with one or more embodiments.

(9) FIG. 8 illustrates a schematic diagram of an intelligent selection and execution platform

utilizing a model selection machine-learning model to select a designated machine-learning model for executing a task in accordance with one or more embodiments.

(10) FIG. 9 illustrates a schematic diagram of an intelligent selection and execution platform adding an additional machine-learning model to a plurality of machine-learning models in accordance with one or more embodiments.

(11) FIGS. 10A-10B illustrate an example diagram of an intelligent selection and execution platform receiving user feedback data and utilizing the user feedback data to update parameters of a model selection machine-learning model in accordance with one or more embodiments.

(12) FIG. 11 illustrates an example series of acts for selecting a designated machine-learning model in accordance with one or more embodiments.

(13) FIG. 12 illustrates an example diagram of an intelligent selection and execution platform selecting a hardware environment for executing a task in accordance with one or more embodiments.

(14) FIG. 13 illustrates a schematic diagram of an intelligent selection and execution platform determining task routing metrics for a plurality of hardware environments in accordance with one or more embodiments.

(15) FIG. 14 illustrates a schematic diagram of an intelligent selection and execution platform utilizing a model hardware allocating machine-learning model to select a designated hardware environment in accordance with one or more embodiments.

(16) FIG. 15 illustrates an example series of acts for selecting a designated hardware environment in accordance with one or more embodiments.

(17) FIG. 16 illustrates a schematic diagram of an intelligent selection and execution platform selecting fallback machine-learning models for executing a task in accordance with one or more embodiments.

(18) FIG. 17 illustrates an intelligent selection and execution platform utilizing a model selection machine-learning model to select a primary machine-learning model and a fallback machine-learning model in accordance with one or more embodiments.

(19) FIG. 18 illustrates an intelligent selection and execution platform utilizing a fallback machine-learning model to execute a task when a primary machine-learning model is unavailable in accordance with one or more embodiments.

(20) FIGS. 19A-19B illustrate an intelligent selection and execution platform utilizing a hardware allocating machine-learning model to select a primary hardware environment and a fallback hardware environment in accordance with one or more embodiments.

(21) FIGS. 20A-20B illustrate an intelligent selection and execution platform detecting a primary hardware environment is unavailable and utilizing a fallback hardware environment to execute a task in accordance with one or more embodiments.

(22) FIG. 21 illustrates an example series of acts for selecting a primary machine-learning model and a fallback machine-learning model in accordance with one or more embodiments.

(23) FIG. 22 illustrates a schematic diagram of an intelligent selection and execution platform intelligently initiating and pausing training tasks based in accordance with one or more embodiments.

(24) FIG. 23 illustrates an example graph depicting bandwidth availability based on usage throughout time in accordance with one or more embodiments.

(25) FIG. 24 illustrates an example series of acts for initiating and pausing training tasks in accordance with one or more embodiments.

(26) FIG. 25 illustrates an example diagram of an intelligent selection and execution platform training a smart pocket machine-learning model in accordance with one or more embodiments.

(27) FIG. 26 illustrates a block diagram of an example computing device for implementing one or more embodiments.

(28) FIG. 27 illustrates a network environment of a content management system in accordance with

one or more embodiments.

DETAILED DESCRIPTION

(29) This disclosure describes one or more embodiments of an intelligent selection and execution platform that assigns or allocates workloads for execution by specific artificial intelligence models and hardware infrastructure. For example, in one or more embodiments, upon receiving a request to initiate a task and determining a hardware load, the intelligent selection and execution platform identifies and designates a machine-learning model for executing the task (or workload) and allocates hardware, such as available GPUs or CPUs, for executing the task (or workload).

Moreover, in one or more embodiments, the intelligent selection and execution platform selects or designates one or more additional models and/or hardware environments as a fallback if the selected model (or hardware) is unavailable for executing the task.

(30) As mentioned, in one or more embodiments, the intelligent selection and execution platform selects a designated model for executing a particular task or workload. In particular, the intelligent selection and execution platform selects a designated machine-learning model from among one or more trained machine-learning models of the intelligent selection and execution platform and various third-party machine-learning models (e.g., models exterior to the intelligent selection and execution platform). For example, upon receiving a request to execute a task, the intelligent selection and execution platform extracts workload features defining characteristics of the task and determines task routing metrics for the various machine-learning models in order to select a designated model as an optimal machine-learning model for executing the task. In some cases, the intelligent selection and execution platform utilizes a model selection machine-learning model to select a designated machine-learning model. In addition, in some instances, the intelligent selection and execution platform utilizes additional information or metrics to select a designated machine-learning model, such as historical quality metrics based on historical user feedback data or software domain analysis information.

(31) As previously mentioned, in one or more embodiments, in addition to selecting a designated machine-learning model, the intelligent selection and execution platform also selects a hardware environment for executing a task. In particular, the intelligent selection and execution platform selects a hardware environment from one or more hardware environments, including hardware environments local to the intelligent selection and execution platform and third-party hardware environments (e.g., external to the intelligent selection and execution platform). For example, upon receiving a request to execute a task, the intelligent selection and execution platform extracts workload features defining characteristics of the task and determines task routing metrics for the various hardware environments in order to select a designated hardware environment for executing a task. In some cases, the intelligent selection and execution platform utilizes a hardware allocating machine-learning model to select a designated hardware environment for executing a task.

Additionally, in some instances, the intelligent selection and execution platform utilizes additional information or metrics for selecting a hardware environment, such as information received from probabilistic load balancing. Moreover, in one or more embodiments, the intelligent selection and execution platform selects a designated hardware environment by selecting multiple hardware environments to execute a task, such as by allocating a first portion of a task to a first hardware environment and a second portion of a task to a second hardware environment.

(32) In addition, as also mentioned, in one or more embodiments, in addition to selecting a model for executing a task, the intelligent selection and execution platform assigns one or more additional models as fallback (or failsafe) models. In particular, the intelligent selection and execution platform assigns additional models as fallback (or failsafe) models to tasks or workloads to prevent execution failures when the selected (or primary) model is unavailable. For example, the intelligent selection and execution platform can select a primary (or designated) machine-learning model and a fallback machine-learning model for executing the task and, thus, if the primary machine-learning model is unavailable (e.g., fails or is busy) for executing the task, the intelligent selection and

execution platform can efficiently defer to the fallback machine-learning model. In one or more embodiments, in addition to (or in lieu of) selecting a primary machine-learning model and a fallback machine-learning model, the intelligent selection and execution platform also selects a primary hardware environment and a fallback hardware environment. Thus, if the primary hardware environment is unavailable (e.g., fails or is busy), the intelligent selection and execution platform can defer to the fallback hardware environment for executing the task.

(33) As briefly mentioned, in one or more embodiments, the intelligent selection and execution platform selects designated machine-learning models and/or hardware environments based in part on the workload features of the task. Specifically, the intelligent selection and execution platform receives workload data requesting the execution of a task and extracts workload features defining characteristics of the task. For example, the intelligent selection and execution platform can extract workload features by determining an estimated processing requirement and an estimated storage requirement for executing the task on each machine-learning model (e.g., the trained machine-learning model and the third-party machine-learning models) running on various hardware environments (e.g., the hardware environment and the third-party hardware environments).

(34) As also briefly mentioned, in one or more embodiments, the intelligent selection and execution platform selects a designated machine-learning model and/or hardware environment based in part on task routing metrics of various machine-learning models and/or hardware environments. In particular, the intelligent selection and execution platform determines task routing metrics by determining various metrics that indicate whether or not a machine-learning model and/or hardware environment are available to execute a task. For example, the intelligent selection and execution platform can determine task routing metrics by determining a model state, a hardware state, a financial cost metric, an execution time metric, a model fit metric, a capability, or a specialty. Moreover, in some cases, the intelligent selection and execution platform utilizes an optimization metric to select a designated machine-learning model as an optimal machine-learning model for executing a task.

(35) In addition, in one or more embodiments, the intelligent selection and execution platform also executes training tasks based on bandwidth availability. In particular, the intelligent selection and execution platform executes training tasks for particular models for their respective tasks using particular models and/or hardware environments and based on the workload data and bandwidth availability. For example, the intelligent selection and execution platform monitors metrics for a hardware environment and/or model and, upon receiving workload data requesting execution of a training task, initiates the task based on bandwidth availability (e.g., at a time when there is expected to be more bandwidth). Upon detecting a change in bandwidth availability (e.g., there is more traffic and less bandwidth to train), the intelligent selection and execution platform can pause the training task. Moreover, in one or more embodiments, in addition to scheduling and pausing training jobs based on bandwidth availability, the intelligent selection and execution platform can also schedule and initiate batch tasks based on bandwidth availability.

(36) The intelligent selection and execution platform provides a variety of technical advantages relative to conventional systems. For example, by dynamically selecting an optimal machine-learning model, the intelligent selection and execution platform improves flexibility relative to conventional systems. Unlike conventional systems that are limited to only one or a small handful of artificial intelligence models, the intelligent selection and execution platform can select an optimal machine-learning model from a plurality of machine-learning models. In particular, the intelligent selection and execution platform selects an optimal machine-learning model for executing a task based on workload features that define characteristics of a task and task routing metrics that indicate how a given machine-learning model would execute the task. Moreover, unlike conventional systems where developers must create systems that utilize only certain artificial intelligence models, a developer need only provide an API call to the intelligent selection and execution platform that requests the execution of a task using a machine-learning model, and

the intelligent selection and execution platform can select an optimal machine-learning model for executing the task.

(37) In addition, the intelligent selection and execution platform maintains flexibility over time relative to conventional systems. In particular, where conventional systems must be completely rebuilt or reworked in order to interface with new artificial intelligence models, the intelligent selection and execution platform can easily add additional (or new) machine-learning models. For example, the intelligent selection and execution platform adds an additional machine-learning model from which to select an optimal model by determining task routing metrics for executing a task on the additional machine-learning model and updating parameters of a smart pocket machine-learning model or a model selection machine-learning model based on the additional model. Moreover, the intelligent selection and execution platform continues to optimize for the new machine-learning model based on user feedback data of implicit and explicit signals regarding the execution of the task by the machine-learning model. Because the model selection machine-learning model receives continuous feedback and uses the feedback to continuously improve selections of the designated machine-learning models, the intelligent selection and execution platform need only provide the model-selection machine-learning model with data (e.g., quality metrics, access to task routing metrics) in order add an additional machine-learning model. Indeed, unlike conventional systems, the intelligent selection and execution platform **102** improves selections after adding the additional machine-learning model without requiring training that is time-intensive and computationally intensive.

(38) In addition, the intelligent selection and execution platform increases efficiency relative to conventional systems by selecting an optimal machine-learning model for a task. Unlike conventional systems that simply move tasks to any available model and often either use a larger model than necessary or use additional bandwidth attempting to execute a large task on a smaller model, the intelligent selection and execution platform can identify an optimal machine-learning model for executing a certain task. By extracting workload features of the task and determining task routing metrics for executing the task on various machine-learning models, the intelligent selection and execution platform can determine from an API call the amount of processing and computational power needed to execute the task and select an optimal machine-learning model that can execute the task. Indeed, the intelligent selection and execution platform can select a machine-learning model that will efficiently execute the task without wasting processing and computing power executing smaller tasks on large models or attempting to execute a large task on a smaller model.

(39) The intelligent selection and execution platform also increases flexibility in selecting hardware environments relative to conventional systems. The intelligent selection and execution platform is aware of the state of multiple hardware environments and uses workload features and task routing metrics to determine an optimal hardware environment for executing the task. Indeed, the intelligent selection and execution platform uses a multi-modal approach to selecting an optimal hardware environment for executing the task. For example, the intelligent selection and execution platform can select a hardware environment with specific capabilities or specialties (e.g., can run a certain model), which will result in a lower financial cost to execute the task. Indeed, the intelligent selection and execution platform can determine between local hardware systems and third-party hardware systems to determine an optimal hardware system for executing the task.

(40) In addition to increasing flexibility when allocating hardware resources, the intelligent selection and execution platform also increases efficiency relative to conventional systems. Specifically, unlike conventional systems that initiate training tasks that result in less bandwidth, the intelligent selection and execution platform can initiate training tasks based on intelligently determining that there is bandwidth availability and pausing the training task based on determining that there was a change in bandwidth availability. Moreover, the intelligent selection and execution platform can determine if a training task requires continuous execution (e.g., is a hot path task) and

begins the training task based on determining that a hardware environment will have sufficient bandwidth availability for the duration of the task, allowing for efficient training scheduling that does not use all the computational and processing power. In addition, unlike conventional systems, the intelligent selection and execution platform can determine a hardware environment to execute a training task based on where a training set is stored, saving additional computational and processing time that conventional systems use to copy training data.

(41) As illustrated by the foregoing discussion, the present disclosure utilizes a variety of terms to describe the features and advantages of the smart topic generation system. Additional details regarding the meaning of such terms are now provided. For example, as used herein, the term “workload data” refers to data, input, or payload that requests the execution of a task and the necessary information to execute the task. In particular, workload data refers to a request to execute a task along with data or other information that indicates the necessary data to execute the task. For example, workload data can refer to data or information that requests specific features or requirements of the task. To illustrate, workload data can refer to an API call received from a device connected to a content management system that requests the execution of a task using a machine-learning model.

(42) Moreover, as used herein, the term “workload features” refers to data, metrics, or other information that indicate requirements for executing a task or generating an outcome. In particular, workload features indicate estimated computational requirements for executing a task. For example, workload features can comprise the necessary computational power from a CPU and/or GPU, time to execute the task, or other specifics necessary for executing a task. To illustrate, workload features can indicate specifics, such as that a task requires a specific machine-learning model or that a task is customer-facing.

(43) In addition, as used herein, the term “machine-learning model” refers to a computer algorithm or a collection of computer algorithms that automatically improve for a particular task through iterative outputs or predictions based on the use of data. For example, a machine-learning model can utilize one or more learning techniques to improve accuracy and/or effectiveness. Example machine-learning models include various types of neural networks, decision trees, support vector machines, linear regression models, and Bayesian networks. In addition, as used herein, the term “trained machine-learning model” refers to a machine-learning model that is local to a content management system. For example, the trained machine-learning model is hosted, located, stored, or executed within a content management system. Moreover, as used herein, the term “third-party machine-learning model” refers to a machine-learning model that is external to a content management system. For example, a third-party machine-learning model is hosted, located, stored, or executed outside of a content management system. Relatedly, as used herein, the term “designated machine-learning model” refers to a machine-learning model selected by a model selection machine-learning model to execute a task.

(44) Relatedly, the term “neural network” refers to a machine-learning model that can be trained and/or tuned based on inputs to determine classifications, scores, or approximate unknown functions. For example, a neural network includes a model of interconnected artificial neurons (e.g., organized in layers) that communicate and learn to approximate complex functions and generate outputs (e.g., content items or smart topic outputs) based on a plurality of inputs provided to the neural network. In some cases, a neural network refers to an algorithm (or set of algorithms) that implements deep learning techniques to model high-level abstractions in data. A neural network can include various layers, such as an input layer, one or more hidden layers, and an output layer that each perform tasks for processing data. For example, a neural network can include a deep neural network, a convolutional neural network, a transformer neural network, a recurrent neural network (e.g., an LSTM), a graph neural network, or a generative adversarial neural network. Upon training, such a neural network may become a machine-learning model.

(45) In addition, as used herein, the term “large language model” refers to a machine-learning

model trained to perform computer tasks to generate or identify content items in response to trigger events (e.g., user interactions, such as text queries and button selections). In particular, a large language model can be a neural network (e.g., a deep neural network or a transformer neural network) with many parameters trained on large quantities of data (e.g., unlabeled text) using a particular learning technique (e.g., self-supervised learning). For example, a large language model can include parameters trained to generate outputs (e.g., smart topic outputs) based on prompts and/or to identify content items based on various contextual data, including graph information from a knowledge graph and/or historical user account behavior. In some cases, a large language model comprises various commercially available models such as, but not limited to, GPT (e.g., GPT 3.5, GPT 4), ChatGPT, Llama (e.g., Llama2-7B, Llama 3), BERT, Claude, Cohere.

(46) As used herein, the term “smart pocket machine-learning model” or “smart pocket ML model” refers to one or more machine-learning models trained or tuned to select from various model options and hardware environment options for executing a task. For example, a smart pocket machine-learning model is trained to select a machine-learning model and/or a hardware environment for executing a task and/or scheduling tasks and training based on workload data, task routing metrics, and hardware usage metrics. In some cases, the smart pocket machine-learning model is a single machine-learning model or algorithm. In other cases, the smart pocket-machine-learning model is a series or ensemble of machine-learning models working together. In one or more embodiments, the smart pocket machine-learning model is a multi-armed bandit.

(47) As used herein, the term “model selection machine-learning model” refers to a machine-learning model that is trained or tuned to select machine-learning models from a plurality of machine-learning models. In particular, a model selection machine-learning model” is trained, tuned, or optimized to select a machine-learning model for executing a task based on workload data and task routing metrics. For example, a model selection machine-learning model can select a machine-learning model from a trained machine-learning model local to a content management system or one or more third-party machine-learning models in their respective network environments. In some cases, the model selection machine-learning model is integrated with or works in series with a smart pocket machine-learning model or a hardware allocating machine-learning model. In other cases, the model selection machine-learning model selects a model separately from the smart pocket machine-learning model.

(48) As used herein, the term “hardware allocating machine-learning model” refers to a machine-learning model that is trained or tuned to select hardware environments from a plurality of hardware environments. In particular, the term hardware allocating machine-learning model refers to a machine-learning model that is trained, tuned, or optimized to select a hardware environment to allocate computing resources to execute a task based on workload features and task routing metrics. For example, a hardware allocating machine-learning model that can select a model from local hardware environments in a content management system and third-party hardware environments.

(49) In addition, as used herein, the term “hardware environment” refers to physical infrastructure and components that provide the computational and storage capacity necessary for executing tasks. In particular, the hardware environment can refer to graphical processing units (GPUs), computational processing units (CPUs), artificial intelligence systems, or other hardware components that can execute various applications, integrations, and computer-executable instructions in order to execute a task. For example, a hardware environment can provide the memory and computation power for running a machine-learning model to execute a task. In some cases, a hardware environment can refer to a hardware environment local to or integrated with a content management system. In other cases, a hardware environment can refer to a third-party hardware environment that is external to a content management system.

(50) Further, as used herein, the term “task routing metrics” refers to various metrics, data, or information relating to whether a machine-learning model or hardware environment is available to

execute a task. In particular, task routing metrics refer to specific metrics regarding whether or not a machine-learning model or hardware environment is capable of executing a task and metrics that relate to costs for executing a task on various machine-learning models and utilizing various hardware environments. For example, task routing metrics can refer to but are not limited to, a financial cost metric, an execution time metric, a model fit metric, a model capability, a model specialty, a hardware capability, or a hardware specialty.

(51) Also, as used herein, the term “hardware usage metrics” refers to metrics that indicate the extent to which resources in a hardware environment are utilized. In particular, the term hardware usage metrics refers to metrics that encompass the measurement and analysis of data transmission rates across components in a hardware environment, such as processors, memory, storage devices, and network interfaces. For example, hardware usage metrics can indicate the bandwidth usage of a hardware environment.

(52) In addition, as used herein, the term “use time period” refers to a defined interval or segment of time. In particular, the term use time period refers to an amount of time with defined start and end points and can describe various durations, including, but not limited to, a number of minutes, hours, days, or weeks. For example, a use time period may signify a phase or timeframe. Relatedly, as used herein, the term “high-use time period” refers to a use time period where various usage metrics indicate that there are a certain number of users or an amount of data flow within a particular system, model, or hardware environment and that the system is above a certain capacity. For example, a high-use time period can refer to a time when usage metrics indicate that the capacity of a system, model, or hardware environment is at a certain capacity, such as through satisfying a threshold or with a score or a metric. In addition, the term “minimum use time period” refers to a use time period where various usage metrics indicate that there are a certain number of users or an amount of data flow within a particular system, model, or hardware environment and that the system is below a certain capacity. For example, a high-use time period can refer to a time when usage metrics indicate that the capacity of a system, model, or hardware environment is below a certain capacity, such as through satisfying (or not satisfying) a threshold or with a score or a metric.

(53) In addition, as used herein, the term “training task bandwidth usage threshold” refers to a level or threshold for initiating or pausing a training task. In particular, a training task bandwidth usage threshold indicates an amount of bandwidth usage of a system or hardware environment that indicates that there is or is not sufficient bandwidth for a training task. For example, if the bandwidth usage does not satisfy a training task bandwidth usage threshold, then the bandwidth usage indicates there is not sufficient bandwidth for a training task, and conversely, if the bandwidth usage satisfies a training task bandwidth usage threshold, then the bandwidth usage indicates there is sufficient bandwidth for a training task. In some embodiments, the training task bandwidth usage threshold could be when the bandwidth usage is above a certain number or percentage (e.g., above 0.65). In other embodiments, the training task initiation threshold could be when a decision tree answers with a “yes” to questions regarding whether there is a certain amount of bandwidth usage.

(54) As used herein, the term “batch task” refers to a task or set of tasks processed as a single unit, typically without user intervention during execution. In particular, the term batch task refers to a set of tasks that are queued together and executed sequentially or in parallel. In some cases, batch task refers to a task that does not require real-time interaction or feedback within a certain time period. For example, a batch task could refer to a group of tasks that are not user-facing in that there is a likelihood they will not be accessed by a client device within a certain time period. To illustrate, a batch task can refer to data processing, report generation, backups, large-scale computations, or syncing tasks.

(55) Further, as used herein, the term “batch task bandwidth usage threshold” refers to a level or threshold for initiating or pausing a batch task. In particular, a batch task bandwidth usage

threshold indicates an amount of bandwidth usage of a system or hardware environment that indicates that there is or is not sufficient bandwidth for a batch task. For example, if the bandwidth usage does not satisfy a batch task bandwidth usage threshold, then the bandwidth usage indicates there is not sufficient bandwidth for a batch task, and conversely, if the bandwidth usage satisfies a batch task bandwidth usage threshold, then the bandwidth usage indicates there is sufficient bandwidth for a batch task. In some embodiments, the batch task bandwidth usage threshold could be when the bandwidth usage is above a certain number or percentage (e.g., above 0.65). In other embodiments, the batch task initiation threshold could be when a decision tree answers with a “yes” to questions regarding whether there is a certain amount of bandwidth usage.

(56) In addition, as used herein, the term “hot path task” refers to a task that requires immediate, uninterrupted processing because they are essential for real-time operation or system stability. In particular, the term hot path task refers to a task wherein pausing or delaying the task could result in performance degradation or service disruption. For example, in some cases, a hot path task refers to a user-facing task where there is a likelihood that a client device will be waiting for a response. To illustrate, a hot path task can refer to, among others, messages, emails, Slack messages, and real-time streaming processing.

(57) Moreover, as used herein, the term “content item” refers to a digital object or a digital file that includes information interpretable by a computing device (e.g., a client device) to present information to a user. A content item can include a file such as a digital text file, a digital image file, a digital audio file, a webpage, a website, a digital video file, a web file, a link, a digital document file, or some other type of file or digital object. A content item can have a particular file type or file format, which may differ for different types of digital content items (e.g., digital documents, digital images, digital videos, or digital audio files). In some cases, a content item can refer to a remotely stored (e.g., cloud-based) item or a link (e.g., a link to a cloud-based item or a web-based content item) and/or a content clip that indicates (or links) a discrete selection or segmented portion of content from a webpage or some other content item or source. A content item can be editable or otherwise modifiable and can also be sharable from one user account (or client device) to another. In some cases, a content item is modifiable by multiple user accounts (or client devices) simultaneously and/or at different times.

(58) In addition, as used herein, the term “network environment” refers to an environment that houses and facilitates the functioning of software and/or hardware components. In particular, the term network environment refers to various components, software, communications, and storage devices for housing and executing a machine-learning model or hardware environment. For example, a network environment can include the collective infrastructure, architecture, and set of protocols that facilitate the functioning of a machine-learning model and/or hardware environment and communication with various other systems or components. To illustrate, a third-party machine-learning model will have a network environment to facilitate the functioning of the third-party machine-learning model and to communicate with the other systems. Similarly, as another illustration, a third-party hardware environment will have a corresponding network environment to facilitate functioning of the third-party hardware environment to facilitate functioning of the third-party hardware environment and to communicate with other systems.

(59) Additional details regarding the intelligent selection and execution platform will now be provided with reference to the figures. For example, FIG. 1 illustrates a block diagram of a system environment for implementing an intelligent selection and execution platform **102** in accordance with one or more embodiments. An overview of the intelligent selection and execution platform **102** is described in relation to FIG. 1. Thereafter, a more detailed description of the components and processes of the intelligent selection and execution platform **102** is provided in relation to the subsequent figures.

(60) As shown, the environment **100** includes server(s) **110**, client device(s) **112**, database **118**, third-party server(s) **120**, and third-party server(s) **124**. Each of the components of the environment

can communicate via network **116**, and network **116** may be any suitable network over which computing devices can communicate. Example networks are discussed in more detail below in relation to FIGS. **26-27**.

(61) As mentioned above, the environment **100** includes client device(s) **112**. The client device(s) **112** can be one of a variety of computing devices, including a smartphone, a tablet, a smart television, a desktop computer, a laptop computer, a virtual reality device, an augmented reality device, or another computing device as described in relation to FIGS. **26-27**. The client device(s) **112** can communicate with the server(s) **110** via network **116**. For example, the client device(s) **112** can receive user input from a user interacting with the client device(s) **112** (e.g., via the client application **114**) to, for instance, select interface elements to interact with a content management system or to select options that initiate execution of a task. In addition, the intelligent selection and execution platform **102** or the server(s) **110** can receive information relating to various interactions with content items and/or user interface elements based on the input received by the client device(s) **112**.

(62) As shown, the client device(s) **112** can include a client application **114**. In particular, the client application **114** may be a web application, a native application installed on the client device(s) **112** (e.g., a mobile application, a desktop application, etc.), or a cloud-based application where all or part of the functionality is performed by the server(s) **110**. Based on instructions from the client application **114**, the client device(s) **112** can present or display information, including a user interface for interacting with (or collaborating regarding) initiating tasks. Using the client application, the client device(s) **112** can perform (or request to perform) various operations, such as executing a task and/or inputting text comprising actions or prompts to generate a specific output.

(63) As illustrated in FIG. **1**, the environment **100** also includes the server(s) **110**. The server(s) **110** may generate, track, store, process, receive, and transmit electronic data, such as results, actions, determinations, responses, computer code, interactions with interface elements, and/or interactions between user accounts or client devices. For example, the server(s) **110** may receive an indication from the client device(s) **112** of a user interaction selecting an option that initiates a task or inputting text comprising actions or prompts to generate a specific output. In addition, the server(s) **110** can transmit data to the client device(s) **112**. Indeed, the server(s) **110** can communicate with the client device(s) **112** to send and/or receive data via network **116**. In some implementations, the server(s) **110** comprise(s) a distributed server where the server(s) **110** include(s) a number of server devices distributed across the network **116** and located in different physical locations. The server(s) **110** can comprise one or more content servers, application servers, container orchestration servers, communication servers, web-hosting servers, machine-learning servers, and other types of servers.

(64) As shown in FIG. **1**, the server(s) **110** can also include the intelligent selection and execution platform **102** as part of the content management system **108**. The content management system **108** can communicate with the client device(s) **112** to perform various functions associated with the client application **114**, such as managing user accounts, initiating tasks, and/or identifying content items. Indeed, content management system **108** can include a network-based smart cloud storage system to manage, store, and maintain content items and related data across numerous user accounts. In some embodiments, the intelligent selection and execution platform **102** and/or the content management system **108** utilize the database **118** to store and access information such as content items, training data sets, or data and/or information related to executing a task.

(65) As also illustrated in FIG. **1**, the content management system **108** can host a trained machine-learning model **104** and a hardware environment **106**. In particular, the content management system **108** can host a trained machine-learning model **104**, and a hardware environment **106** local to (e.g., a part of or integrated within) the content management system **108**. For example, the content management system **108** utilizes the trained machine-learning model **104** and hardware environment **106** to execute tasks locally within the content management system **108**. Indeed, the content management system **108** trains, maintains, and manages the workload of the trained

machine-learning model **104** and the hardware environment **106**.

(66) As further illustrated in FIG. 1, the environment **100** includes the third-party server(s) **120** that host the third-party machine-learning model(s) **122**. In particular, the third-party machine-learning model(s) **122** communicates with the server(s) **110**, the client device(s) **112**, the database **118**, and/or the third-party server(s) **124** (or the third-party hardware environment(s) hosted on the third-party server(s) **124**) for the intelligent selection and execution platform **102** to select a model and/or a hardware environment for executing a task or to schedule model training. For example, the intelligent selection and execution platform **102** provides domain-specific language segments to the third-party machine-learning model(s) **122**, where the domain-specific language segments indicate data for generating results for various subcomponents. Indeed, the third-party machine-learning model(s) **122** can include a machine-learning model powered by neural networks or other machine-learning architectures for generating responses to text queries. In some cases, the third-party machine-learning model(s) **122** can refer to various third-party machine-learning models (e.g., ChatGPT, Lambda, Llama, BERT, RoBERTa, Turing-NLG, T5, XLNet).

(67) As also illustrated in FIG. 1, the environment **100** includes the third-party server(s) **124** that host the third-party hardware environment(s) **126**. In particular, the third-party hardware environment(s) **126** communicate with the server(s) **110**, the client device(s) **112**, the database **118**, and/or the third-party server(s) **120** to execute tasks, such as by providing resources to execute tasks with a machine-learning model or to train a machine-learning model (e.g., based on training scheduled by the intelligent selection and execution platform **102**). For example, the third-party hardware environment(s) provide computational power and memory for a machine-learning model to execute a task. Indeed, the third-party hardware environment(s) **126** can include graphical processing units and central processing units for executing a task and/or training a machine-learning model. In some cases, the third-party hardware environment(s) **126** can refer to various hardware environment(s) that include infrastructure available for a fee.

(68) In some implementations, though not illustrated in FIG. 1, the environment **100** may have a different arrangement of components and/or may have a different number or set of components altogether. For example, the client device(s) **112** may communicate directly with the intelligent selection and execution platform **102** and not through network **116**. The environment may also include one or more third-party systems, each corresponding to a different data source. In addition, the environment can include the database **118** located external to the server(s) **110** (e.g., in communication via the network **116**) or located on the server(s) **110** and/or on the client device(s) **112**.

(69) As mentioned, the intelligent selection and execution platform **102** can intelligently assign or allocate tasks for execution by specific artificial intelligence machine-learning models and hardware infrastructure. In particular, the intelligent selection and execution platform **102** can utilize a smart pocket ML model that can select an artificial intelligence model, allocate hardware, select fallback machine-learning models and/or hardware, and schedule model training. FIG. 2 illustrates an example diagram of an overview of an intelligent selection and execution platform **102** utilizing a smart pocket ML model in accordance with one or more embodiments.

(70) As illustrated in FIG. 2, the intelligent selection and execution platform **102** includes an API layer **202** that receives prompts or workload data. In particular, the intelligent selection and execution platform **102** receives workload data that indicates (or requests) a desired outcome from a machine-learning model. For example, API layer **202** acts as a central repository for requests from third-party systems or applications to utilize a machine-learning model to execute a task. Moreover, in some instances, requests received by API layer **202** also comprise requests to access or utilize content items stored by the content management system **108** to execute a task. Indeed, API layer **202** receives prompts in a designated format that allows for various other systems to flexibly utilize machine-learning model capabilities by simply sending prompts to the API layer in the designated format. Additional detail regarding the API layer receiving prompts and/or workload

data is provided in relation to FIG. 4 below.

(71) As also illustrated in FIG. 2, the intelligent selection and execution platform **102** includes a smart pocket ML model **204**. In particular, the smart pocket ML model **204** includes various layers for scheduling, routing, and intelligently selecting and designating hardware and software components for executing a particular task based on the availability of hardware and software components. For example, the smart pocket ML model **204** can include a trained model (or heuristic model), such as a neural network or a multi-armed bandit model, to determine how to allocate computing resources and hardware within the intelligent selection and execution platform **102**, the content management system **108**, third-party machine-learning models (e.g., third-party machine-learning model(s) **122**) and/or third-party hardware (e.g., third-party hardware environment(s) **126**).

(72) In one or more embodiments, smart pocket ML model **204** or API layer **202** also include metrics of portions of the intelligent selection and execution platform **102** and/or the content management system **108**. In particular, smart pocket ML model **204** or API layer **202** monitor the status of machine-learning models and/or hardware environments of the intelligent selection and execution platform **102** and/or the content management system **108**. For example, smart pocket ML model **204** or API layer **202** may monitor or have access to metrics and/or data about graphical processing units (GPUs) of the intelligent selection and execution platform **102** and/or the content management system **108** that the intelligent selection and execution platform **102** can utilize in order to select models and/or hardware to execute a particular task.

(73) As illustrated in FIG. 2, the smart pocket ML model **204** can include model selection **206**. In particular, model selection **206** includes selecting a designated machine-learning model for executing a task. Specifically, the smart pocket ML model **204** selects a designated machine-learning model based on the workload features of the task and task routing metrics for a plurality of machine-learning models hosted in various network environments. For example, the plurality of machine-learning models includes a trained machine-learning model within the intelligent selection and execution platform **102** (or the content management system **108**) and third-party machine-learning models hosted on third-party servers. Additional detail regarding the intelligent selection and execution platform **102** selecting and utilizing a machine-learning model to execute a task will be provided in relation to FIGS. 5-10 below.

(74) In addition, as illustrated in FIG. 2, the smart pocket ML model **204** can include hardware allocation **208**. In particular, the smart pocket ML model **204** can select one or more hardware environments to allocate resources for executing a task on a designated machine-learning model. For example, the smart pocket ML model **204** selects a designated hardware environment based on the workload features of the task and task routing metrics for a plurality of hardware environments hosted in various network environments. In some cases, the smart pocket ML model **204** selects a designated hardware environment from a hardware environment within the intelligent selection and execution platform **102** and third-party machine-learning models hosted on third-party servers. More detail regarding the intelligent selection and execution platform **102** selecting a hardware environment to execute a task will be provided in relation to FIGS. 12-15 below.

(75) Moreover, as illustrated in FIG. 2, the smart pocket ML model **204** can include model/hardware fallback **210**. In particular, the smart pocket ML model **204** can select a primary machine-learning model and a fallback machine-learning model for executing a task. In particular, the smart pocket ML model **204** can select a primary machine-learning model as a designated machine-learning model and, if the primary machine-learning model is unavailable, can reassign executing the task to the fallback machine-learning model. For example, the smart pocket ML model **204** extracts workload features defining characteristics of the task, determines task routing metrics for various hardware environments, and selects the primary machine-learning model and the fallback machine-learning model based on the workload features and task routing metrics. More details regarding the intelligent selection and execution platform **102** selecting a primary machine-

learning model and a fallback machine-learning model are provided in relation to FIGS. **16-18** below.

(76) In addition to selecting a primary machine-learning model and a fallback machine-learning model, the smart pocket ML model **204** can select a primary hardware environment and a fallback hardware environment. Specifically, if the primary hardware environment is unavailable, the smart pocket ML model **204** can allocate resources from the secondary hardware environment for executing the task. For example, the smart pocket ML model **204** extracts workload features defining characteristics of the task, determines task routing metrics for a plurality of hardware environments hosted in various network environments, and selects a primary hardware environment and a fallback hardware environment based on the workload features and the task routing metrics. Additional details regarding the intelligent selection and execution platform **102** selecting a primary hardware environment and a fallback hardware environment are provided with relation to FIG. **16** and FIGS. **19A-20B** below.

(77) Further, as illustrated in FIG. **2**, the smart pocket ML model **204** can include model training initiating **212**. In particular, the smart pocket ML model **204** initiates training tasks based on bandwidth availability and pauses training tasks based on bandwidth availability. For example, the smart pocket ML model **204** monitors hardware usage metrics for hardware environments in various network environments and based on the workload data of the training task and bandwidth availability indicated by the hardware usage metrics, the smart pocket ML model **204** can initiate the training task. Moreover, if the smart pocket ML model **204** detects a change in bandwidth availability, the smart pocket ML model **204** can pause the training task. Additional detail regarding the intelligent selection and execution platform **102** scheduling model training will be provided in relation to FIGS. **22-24** below.

(78) As shown in FIG. **2**, the intelligent selection and execution platform **102** can execute task **214**. In particular, the intelligent selection and execution platform **102** can execute the task based on determinations or decisions from the smart pocket ML model **204**. For example, the intelligent selection and execution platform **102** can execute the task using the designated machine-learning model and/or designated hardware environment or on the fallback machine-learning model and/or fallback hardware environment. As another example, the intelligent selection and execution platform **102** can execute the task by initiating and pausing training tasks based on bandwidth availability.

(79) As previously mentioned, the intelligent selection and execution platform **102** utilizes a smart pocket ML model to select machine-learning models, allocate hardware for tasks, and schedule model training. In particular, the intelligent selection and execution platform **102** utilizes a smart pocket machine-learning model to select models and/or hardware environments for executing a task and/or scheduling model training. FIG. **3** illustrates an example diagram of an overview of an intelligent selection and execution platform utilizing a smart pocket machine-learning model in accordance with one or more embodiments.

(80) As shown in FIG. **3**, the intelligent selection and execution platform **102** utilizes the smart pocket machine-learning model **308** to generate output **314**. In particular, the smart pocket machine-learning model **308** generates output **314** based on characteristics of a task and various metrics that indicate costs for executing the task or whether a machine-learning model and/or hardware environment is available to execute a task. For example, the intelligent selection and execution platform **102** provides characteristics of the task to smart pocket machine-learning model **308** by extracting workload features **302** from workload data. In some cases, workload features **302** indicate estimated processing requirements or estimated storage requirements for executing a task. Moreover, task routing metrics **304** indicate whether a machine-learning model and/or hardware environment is available to execute a task and/or costs for executing the task on the machine-learning model and/or hardware environment. In addition, the intelligent selection and execution platform **102** can identify or determine characteristics **306** that indicate certain characteristics

and/or specialties of a machine-learning model or hardware environment. Example task routing metrics for machine-learning models are discussed further in relation to FIG. 7 below. In addition, example task routing metrics for hardware environments are discussed further in relation to FIG. 12 below.

(81) In one or more embodiments, smart pocket machine-learning model **308** is trained, tuned, or optimized to generate output **314**. In some cases, smart pocket machine-learning model **308** is a multi-armed bandit model or optimization model that is optimized to generate output **314**, such as a designated (and fallback) machine-learning model from a plurality of machine-learning models, a designated (and fallback) hardware environment from a plurality of hardware environments or initiating model training. In other cases, smart pocket machine-learning model **308** is a machine-learning model or neural network that is trained or tuned to select a designated (and fallback) machine-learning model from a plurality of machine-learning models, select a designated (and fallback) hardware environment from a plurality of hardware environments, or schedule model training.

(82) As indicated in FIG. 3, in one or more embodiments, smart pocket machine-learning model **308** comprises multiple machine-learning models working in series to generate output **314**. Specifically, smart pocket machine-learning model **308** can utilize a model selection machine-learning model **310** and a hardware allocating machine-learning model **312**. For example, the intelligent selection and execution platform **102** utilizes the model selection machine-learning model **310** to select a designated machine-learning model and additional machine-learning models as fallback machine-learning models for executing a task. Moreover, the intelligent selection and execution platform **102** utilizes the hardware allocating machine-learning model **312** to select a designated hardware environment and additional hardware environments for executing the task. Additional detail regarding the intelligent selection and execution platform **102** utilizing a model selection machine-learning model is provided below in relation to FIG. 8 below. Additional details regarding the intelligent selection and execution platform **102** utilizing a hardware allocating machine-learning model are provided below, in relation to FIG. 14 below.

(83) In one or more embodiments, the intelligent selection and execution platform **102** utilizes a computing stack to execute various processes or actions. In particular, the intelligent selection and execution platform **102** utilizes various layers of the computing stack to perform various actions, such as selecting machine-learning models and/or hardware environments, selecting primary and fallback machine-learning models, and scheduling and/or initiating training tasks. FIG. 4 illustrates an example diagram of an overview of layers of an intelligent selection and execution platform in accordance with one or more embodiments.

(84) As illustrated in FIG. 4, the intelligent selection and execution platform **102** comprises an API layer **402** that receives workload data. In particular, API layer **402** is an application programming interface (API) layer that receives workload data that requests the execution of a task using a machine-learning model. For example, the workload data can comprise data or information that represents the level and nature of the task or estimated hardware or software requirements for executing the task. In some cases, workload data comprises parameters for executing the task, such as a request to use (or not use) a particular machine-learning model or hardware environment.

(85) In one or more embodiments, API layer **402** receives workload data from a client device connected by a network to a content management system **108**. Specifically, through API layer **402**, other systems, applications, or devices can connect to resources of the content management system **108** or utilize resources of the content management system **108** when executing the task. For example, workload data may comprise a request to execute a task using a machine-learning model and content items within the content management system **108**. Additional details regarding the intelligent selection and execution platform **102** receiving workload data from a client device are provided with relation to FIG. 6 below.

(86) In one or more embodiments, workload data is in a designated format for the API layer **402**.

Specifically, the intelligent selection and execution platform **102** requires a designated format for third-party systems to access the intelligent selection and execution platform **102** (or the content management system **108**) through API layer **402**. For example, workload data provided to API layer **402** can comprise a request to execute a task using a machine-learning model. Moreover, workload data provided to API layer **402** does not need to specify a machine-learning model on which to execute the task. Indeed, so long as a request to API layer **402** is in the designated format, the intelligent selection and execution platform **102** can select a designated (optimal) machine-learning model and/or hardware environment for executing the task.

(87) Moreover, in one or more embodiments, API layer **402** is part of the content management system **108** and is not limited to the intelligent selection and execution platform **102**. Specifically, API layer **402** can be an API layer of the content management system **108**, where API layer **402** is a central repository for third-party systems and applications that would like to access the content management system **108**. Upon receiving an API call in API layer **402**, the content management system **108** can utilize the intelligent selection and execution platform **102** to select a designated (and fallback) machine-learning model and/or designated hardware environment for executing a task. For example, when the content management system **108** receives an API call at API layer **402** and identifies that workload data included in the API call comprises a request to execute a task using a machine-learning model, the content management system **108** can utilize the intelligent selection and execution platform **102** to select a designated (and fallback) machine-learning model and/or designated (and fallback) hardware environment for executing the task.

(88) In addition to receiving workload data, API layer **402** monitors hardware usage metrics. Specifically, the API layer **402** monitors hardware usage metrics for a plurality of hardware environments in respective network environments and acts as a central location with knowledge of availability for the hardware environments. For example, the API layer **402** may monitor hardware usage metrics for various GPU and/or CPU units of the plurality of hardware environments.

(89) As further illustrated in FIG. 4, the intelligent selection and execution platform **102** comprises a prompt layer **404**. In one or more embodiments, when the request to initiate a task indicates or requests the intelligent selection and execution platform **102** to execute a task using a large language model as the machine-learning model, prompt layer **404** provides prompts for a large language model to generate a certain output. In some cases, prompt layer **404** utilizes prompt templates that the intelligent selection and execution platform **102** can utilize to generate a certain output from a large language model. For example, each prompt template comprises a previously written and tested prompt that generates a consistent output from a large language model.

(90) In other cases, prompt layer **404** receives prompts from prompt generator **420**. In particular, prompt generator **420** is a paradigm that utilizes algorithms to generate prompts for large language models by optimizing large language model prompts and weights. For example, prompt generator **420** generates and optimizes prompts for large language models based on a training data set and corresponding metrics. In addition, the prompt generator **420** can fine-tune the large language model, resulting in a smaller model optimized for a specific task. Prompt layer **404** can also store prompts generated by prompt generator **420** for future use by the intelligent selection and execution platform **102**.

(91) In some embodiments, the intelligent selection and execution platform **102** may receive a prompt as part of workload data received in API layer **402**. Specifically, the intelligent selection and execution platform **102** can receive a prompt from a third-party system or application as part of the workload data of an API call. In some cases, the intelligent selection and execution platform **102** can provide the prompt directly from the API layer to the prompt layer (e.g., without utilizing the prompt generator **420**). In other cases, the intelligent selection and execution platform **102** will utilize the prompt generator **420** to optimize a prompt received from a third-party application or third-party system.

(92) As also illustrated in FIG. 4, the intelligent selection and execution platform **102** includes a

model and hardware selection layer **406**. In particular, model and hardware selection layer **406** selects a designated (and fallback) machine-learning model for executing a task and a designated (and fallback) hardware environment for executing a task. In some cases, the model and hardware selection layer **406** comprises one or more machine-learning models, such as the smart pocket machine-learning model **308**, the model selection machine-learning model **310**, or the hardware allocating machine-learning model **312**.

(93) As illustrated in FIG. **4**, the intelligent selection and execution platform **102** comprises a leaderboard **418**. In one or more embodiments, leaderboard **418** is a database that holds performance metrics (or task routing metrics), including latency metrics and model quality metrics, for all machine-learning models, both a trained machine-learning model local to the content management system **108** and third-party machine-learning models. For example, leaderboard **418** comprises leaderboards grouped by task, with each task featuring a list of machine-learning models and their associated performance metrics. Moreover, leaderboard **418** can also comprise one or more leaderboards for specific tasks. For example, leaderboard **418** can comprise a summarization leaderboard that collects all machine-learning models that support summarization and the corresponding metrics for the machine-learning models.

(94) Moreover, as illustrated in FIG. **4**, model and hardware selection layer **406** can read data and other information from a leaderboard **418**. model and hardware selection layer **406** can utilize data, metrics, and other information from leaderboard **418** to select designated (and fallback) machine-learning models and designated (and fallback) hardware environments. For example, leaderboard **418** comprises model quality metrics based on previous output generated by each machine-learning model. The intelligent selection and execution platform **102** generates model quality metrics based on user feedback data from evaluation and benchmarking pipeline **416** indicating user satisfaction with the performance of corresponding machine-learning models.

(95) As further illustrated in FIG. **4**, the intelligent selection and execution platform **102** comprises a security and moderation layer **408**. In particular, security and moderation layer **408** provides security and moderation for prompts for machine-learning models to protect against breaches or other security concerns, particularly with open-source machine-learning models. For example, the intelligent selection and execution platform **102** can utilize third-party security systems that specialize in providing security protection for artificial intelligence models (e.g., Lakera Guard or Purple Llama).

(96) In addition, as illustrated in FIG. **4**, the intelligent selection and execution platform **102** comprises a routing layer **410**. Specifically, the routing layer **410** routes data or other information between components of the intelligent selection and execution platform **102**. For example, the routing layer **410** routes data between machine-learning models, hardware environments, and the content management system, among others, in order to determine task routing metrics, receive workload data comprising a request to execute a task, and receive prompts.

(97) As further illustrated in FIG. **4**, the intelligent selection and execution platform **102** comprises a model dispatch layer **412** that provides data and other information to selected machine-learning models. Specifically, model dispatch layer **412** provides prompts to designated machine-learning models to execute a task. Moreover, model dispatch layer **412** can also receive user feedback data indicating user satisfaction with the execution of the prompt and provide the user feedback data to the evaluation and benchmarking pipeline **416**.

(98) As also illustrated in FIG. **4**, the intelligent selection and execution platform **102** comprises a hardware dispatch layer **414**. In particular, the hardware dispatch layer **414** provides data to (and communicates with) hardware environments for executing tasks using selected models. For example, hardware dispatch later can be a hardware environment of the intelligent selection and execution platform **102** or a third-party hardware environment. In some cases, the hardware environment is Nvidia, AWS, Google cloud, or other suitable third-party hardware environment.

(99) Moreover, as illustrated in FIG. **4**, the intelligent selection and execution platform **102** can

comprise annotated data **422**. In particular, annotated data **422** comprises labeled output from various machine-learning models. In one or more embodiments, the intelligent selection and execution platform **102** accesses annotated data **422** to evaluate quality for machine-learning models based on performance metrics, financial cost metrics, or execution time metrics. In some cases, the intelligent selection and execution platform **102** will use annotated data **422** to calibrate model ranking in leaderboard **418** and to train and optimize prompts for prompt generator **420**.

(100) As previously mentioned, the intelligent selection and execution platform **102** selects a model for executing a task. In particular, the intelligent selection and execution platform **102** selects a model for the task based on workload features defining characteristics of the task and task routing metrics of machine-learning models. FIG. 5 illustrates a schematic diagram of an intelligent selection and execution platform selecting a designated model for executing a task in accordance with one or more embodiments.

(101) As shown in FIG. 5, the intelligent selection and execution platform **102** performs an act **502** and receives workload data requesting the execution of a task. In particular, the intelligent selection and execution platform **102** receives workload data from a device connected to the content management system that requests the execution of a task. For example, workload data comprises the necessary data and/or information for executing the task, such as a request to execute the task using a machine-learning model. In one or more embodiments, the intelligent selection and execution platform **102** receives workload data through an API layer, as described in relation to FIG. 4 above.

(102) In one or more embodiments, the intelligent selection and execution platform **102** workload data comprises parameters for executing a task. Specifically, the workload data comprises parameters comprising requests or specifications for executing a task. For example, workload data can comprise a parameter that instructs the intelligent selection and execution platform **102** to utilize a certain machine-learning model for executing a task or to execute the task using a certain hardware environment. As another example, workload data can comprise a constraint that instructs the intelligent selection and execution platform **102** not to utilize a certain machine-learning model or hardware environment to execute a task. Additional details regarding the intelligent selection and execution platform **102** receiving workload data and parameters or constraints for executing a task are provided with relation to FIG. 6 below.

(103) As also shown in FIG. 5, the intelligent selection and execution platform **102** performs an act **504** and extracts, from the workload data, workload features defining characteristics of the task. In particular, the intelligent selection and execution platform **102** extracts workload features by determining an amount of processing and storage needed for executing the task. For example, the intelligent selection and execution platform **102** can compute the computational cost for executing a task using a specific machine-learning model. In some cases, the intelligent selection and execution platform **102** determines processing requirements and storage requirements for executing the task on trained machine-learning models within the content management system and third-party machine-learning models. Additional details regarding extracting workload features from workload data are discussed in FIG. 6 below.

(104) As also shown, the intelligent selection and execution platform **102** performs an act **506** and determines task routing metrics for a plurality of models hosted in respective network environments. In particular, the intelligent selection and execution platform **102** determines task routing metrics that indicate the availability of machine-learning models, such as whether they are available to perform the task. For example, the intelligent selection and execution platform **102** determines a model state that indicates a current operation status, resource usage metrics, or operational metrics for each machine-learning model of the plurality of machine-learning models. As another example, the intelligent selection and execution platform **102** determines a model capability indicating that a machine-learning model is able to execute a certain task or perform a certain action, a model specialty indicating that a machine-learning model is adept at performing a

certain action or task, or a model fit indicating alignment of the machine-learning model with a task.

(105) In addition, in one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics for executing the task on various machine-learning models. In particular, the intelligent selection and execution platform **102** determines task routing metrics indicating costs or lengths of executing a task on each machine-learning model of the plurality of machine-learning models. For example, the intelligent selection and execution platform **102** determines financial cost metrics, execution time metrics, or execution cost metrics for executing the task on each machine-learning model in the plurality of machine-learning models. Additional details regarding the intelligent selection and execution platform **102** determining task routing metrics for a plurality of machine-learning models hosted in respective network environments are provided in relation to FIG. 7 below.

(106) As illustrated in FIG. 5, the intelligent selection and execution platform **102** can also perform an act **508** and select a designated machine-learning model for executing the task. In particular, the intelligent selection and execution platform **102** uses workload features and task routing metrics to select a designated machine-learning model for executing the task. For example, the intelligent selection and execution platform **102** selects a machine-learning model by using a trained model (or heuristic model), such as a neural network or a multi-armed bandit model, to determine an optimal model for a task. In one or more embodiments, the intelligent selection and execution platform **102** utilizes a model selection machine-learning model to select a designated machine-learning model. In some cases, the model selection machine-learning model is trained, tuned, or optimized to select machine-learning models for executing tasks.

(107) As mentioned, in one or more embodiments, the intelligent selection and execution platform **102** selects a designated machine-learning model for executing a task based on workload features and task routing metrics. Specifically, the intelligent selection and execution platform **102** utilizes the model selection machine-learning model to analyze the workload features and task routing metrics and select a designated machine-learning model for executing the task. For example, the model selection machine-learning model generates an optimization metric based on the workload features and the task routing metrics and selects the designated machine-learning model based on the optimization metric. Additional detail regarding the intelligent selection and execution platform **102** utilizing a model selection machine-learning model to select a designated machine-learning model for executing a task is discussed further in relation to FIG. 8 below.

(108) In one or more embodiments, the intelligent selection and execution platform **102** can add additional machine-learning models to the plurality of machine-learning models. Specifically, the intelligent selection and execution platform **102** updates parameters of the smart pocket machine-learning model or the model selection machine-learning model based on task routing metrics of the additional machine-learning model. For example, upon updating the model selection machine-learning model, the intelligent selection and execution platform **102** can select the additional trained model from among the plurality of trained machine-learning models. Additional detail regarding the intelligent selection and execution platform **102** adding an additional machine-learning model to a plurality of machine-learning models is provided below in relation to FIG. 9 below.

(109) Further, in one or more embodiments, after selecting the model, the intelligent selection and execution platform **102** trains the model (or heuristic model) used to select from the various models based on use feedback. In particular, the intelligent selection and execution platform **102** can update the parameters of the model-selection machine-learning model based on user feedback data upon execution of the task. For example, the intelligent selection and execution platform **102** receives user feedback data by receiving explicit feedback (e.g., selecting quality options) or implicit feedback (e.g., user reactions to the executed task). Additional details regarding receiving user feedback data and using user feedback data to update parameters of the model-selection

machine-learning model are provided in relation to FIGS. **10A-10B** below.

(110) As previously mentioned, the intelligent selection and execution platform **102** extracts workload features from workload data. In particular, the intelligent selection and execution platform **102** receives, from a client device, workload data requesting the execution of a task and extracts workload features defining characteristics of the task from the workload data. FIG. **6** illustrates a schematic diagram of an intelligent selection and execution platform **102** extracting workload features from workload data requesting execution of a task in accordance with one or more embodiments.

(111) As illustrated in FIG. **6**, the intelligent selection and execution platform **102** receives workload data **604** from a client device **602**. Specifically, the intelligent selection and execution platform **102** receives workload data **604** from client device **602** that is connected to a content management system (e.g., content management system **108**). In one or more embodiments, the intelligent selection and execution platform **102** receives workload data **604** from an application associated with the content management system **108**. In some cases, the intelligent selection and execution platform **102** receives workload data **604** based on user input within an application (e.g., within a graphical user interface) of content management system **108** on client device **602**. In other cases, the intelligent selection and execution platform **102** receives workload data **604** based on user input within a third-party application on client device **602** that is connected to the content management system **108**.

(112) As previously mentioned in relation to FIG. **4** above, in one or more embodiments, the intelligent selection and execution platform **102** receives workload data **604** through an application programming interface (API). Specifically, the intelligent selection and execution platform **102** communicates with third-party applications on the client device **602** through an API (e.g., API layer **402**) to exchange data, request execution of tasks, or perform additional actions. For example, the intelligent selection and execution platform **102** receives workload data **604** as part of an API call in a designated format required by the intelligent selection and execution platform **102** (or the content management system **108**) and where the API call initiated from user interactions with the third-party application on the client device.

(113) As previously mentioned, in some embodiments, workload data **604** comprises a request for the execution of a task. In particular, the request for execution of a task comprises a request to execute the task using a machine-learning model. Indeed, workload data **604** needs simply request the execution of a task using a machine-learning model for the intelligent selection and execution platform **102** to select a designated (and fallback) machine-learning model and/or designated (and fallback) hardware environment for executing the task. For example, a third-party application or system may provide workload data requesting the execution of a task using a machine-learning model (e.g., a large language model) as part of an API call, and as long as the API call is in the designated format, the intelligent selection and execution platform **102** can select designated machine-learning models and/or designated hardware environments for executing the task.

(114) In addition, in one or more embodiments, workload data **604** comprises parameters for executing the task. Specifically, parameters indicate a condition for executing the task that indicates a desire or request for the intelligent selection and execution platform **102** to utilize when selecting a designated machine-learning model and/or designated hardware environment. In one or more embodiments, a parameter can indicate a request to execute a task using a preferred machine-learning model or a preferred hardware environment. Specifically, a parameter that indicates a preferred machine-learning model or a preferred hardware environment will also specify a model identification that indicates a preferred machine-learning model and/or preferred hardware environment. For example, a preferred machine-learning model can be a personalized machine-learning model that is trained or tuned to provide specific output.

(115) Moreover, in one or more embodiments, parameters are not included in workload data **604** but are provided or indicated in another portion of the intelligent selection and execution platform

102. For example, in cases where a parameter is indicated by a service agreement with a third-party system or application, the intelligent selection and execution platform **102** may identify the parameters based on other data. For instance, the intelligent selection and execution platform **102** may train a model selection machine-learning model to identify that workload data **604** from certain third-party systems or applications comprise certain parameters (e.g., a certain machine-learning model or certain hardware environment).

(116) In one or more embodiments, a parameter can indicate a constraint for executing a task. Specifically, a constraint indicates a desire or request for the intelligent selection and execution platform **102** to not select certain machine-learning models or hardware environments. For example, a parameter may comprise a constraint to not use any third-party machine-learning models or hardware environments and use only the trained machine-learning model and hardware environment of the intelligent selection and execution platform **102**. As another example, a parameter can comprise a constraint to refrain from using any personalized machine-learning models and use only generic machine-learning models.

(117) Further, in some embodiments, a parameter may indicate whether or not to use another machine-learning model or hardware environment if the preferred machine-learning model or the preferred hardware environment is unavailable. For example, in some cases, a parameter may indicate that another machine-learning model or hardware environment may be selected if the preferred machine-learning model or preferred hardware environment is unavailable. In other cases, a parameter may indicate selecting only the preferred machine-learning model or preferred hardware environment and performing another action (e.g., providing an error message) if the preferred machine-learning model or the preferred hardware environment is unavailable.

(118) As shown in FIG. **6**, workload data **604** can optionally comprise a prompt **606**. Specifically, when workload data **604** indicates that the task should utilize a large language model to execute the task, workload data **604** can include prompt **606** that comprises instructions for a large language model to generate a specific output when (or as part of) executing the task. For example, a third-party application or system may provide the prompt as part of workload data **604**. The intelligent selection and execution platform **102** may utilize prompt **606** as provided in relation to workload data **604** or may optimize the prompt (e.g., using prompt generator **420**).

(119) As previously mentioned, the intelligent selection and execution platform **102** extracts workload features from workload data. As illustrated in FIG. **6**, the intelligent selection and execution platform **102** extracts workload features **608** from workload data **604**. In one or more embodiments, the intelligent selection and execution platform **102** extracts workload features **608** from workload data **604** by determining estimated processing requirements **610** for executing the task requested in workload data **604**. Specifically, the intelligent selection and execution platform **102** generates estimated processing requirements by determining an estimated processing power requirement, an estimated memory requirement, or an estimated hardware environment requirement. For example, an estimated memory requirement indicates an amount of RAM required to execute the task, and an estimated hardware environment requirement indicates the estimated CPU or GPU requirements for executing the task.

(120) As also illustrated in FIG. **6**, in some embodiments, the intelligent selection and execution platform **102** extracts workload features **608** by determining estimated storage requirements **612**. In particular, the intelligent selection and execution platform **102** determines estimated storage requirements by determining data storage device requirements for executing the task. For example, in one or more embodiments, the intelligent selection and execution platform **102** determines an amount of data storage requirements or a type of data storage requirements for executing the task.

(121) As previously mentioned, the intelligent selection and execution platform **102** determines task routing metrics for machine-learning models. In particular, the intelligent selection and execution platform **102** determines task routing metrics for a trained machine-learning model of the intelligent selection and execution platform **102** and one or more third-party trained models. FIG. **7**

illustrates a schematic diagram of an intelligent selection and execution platform determining task routing metrics for a plurality of machine-learning models hosted in respective network environments in accordance with one or more embodiments.

(122) As illustrated in FIG. 7, in some embodiments, the intelligent selection and execution platform **102** determines task routing metrics **708** for a plurality of machine-learning models. Specifically, the intelligent selection and execution platform **102** determines task routing metrics **708** that indicate costs, timeframes, capabilities, fit, or feedback for each model of a plurality of machine-learning models and which the intelligent selection and execution platform **102** can further provide as input for a model selection machine-learning model. As illustrated, the intelligent selection and execution platform **102** can determine task routing metrics **708** for machine-learning models in their respective environments, including a trained machine-learning model **702** that resides on (or is local to) the intelligent selection and execution platform **102** and third-party machine-learning model(s) **706** that reside on third-party server(s) **704**. Indeed, the intelligent selection and execution platform **102** determines task routing metrics for the trained machine-learning model **702** or the third-party machine-learning model(s) **706** in order to determine an optimal machine-learning model (e.g., based on an optimization metric).

(123) In one or more embodiments, as shown in FIG. 7, the intelligent selection and execution platform **102** can determine task routing metrics **708** by determining financial cost metrics **710**. In particular, intelligent selection and execution platform **102** can determine financial cost metrics **710** by determining financial costs for executing the task on the trained machine-learning model **702** or the third-party machine-learning model(s) **706** (e.g., how much it will cost to execute the task using that model). For example, the intelligent selection and execution platform **102** can determine a currency amount for executing a task on the trained machine-learning model and the third-party machine-learning model(s). In one or more embodiments, the intelligent selection and execution platform **102** can also determine financial cost metrics **710** based on a use time period during which the intelligent selection and execution platform **102** executes the task. For example, financial cost metrics **710** may be higher when executing a task on a third-party machine-learning model at a high-use (or peak) time period than when executing the task during a minimum-use (or non-peak) time period.

(124) As also shown in FIG. 7, in some embodiments, the intelligent selection and execution platform **102** can determine task routing metrics **708** by determining execution time metrics **712**. Specifically, the intelligent selection and execution platform **102** determines execution time metrics **712** by determining an execution time for executing the task on the trained machine-learning model **702** or the third-party machine-learning model(s) **706**. For example, the intelligent selection and execution platform **102** can determine that a first machine-learning model will execute a task in a first time frame and a second machine-learning model will execute the task in a second time frame based on execution time metrics **712** for each respective model. To illustrate, a third-party machine-learning model(s) **706** could execute the task in less time than a trained machine-learning model **702** (or vice versa).

(125) In some embodiments, as shown in FIG. 7, the intelligent selection and execution platform **102** determines task routing metrics **708** by determining execution cost metrics **714**. In particular, the intelligent selection and execution platform **102** determines execution cost metrics **714** by determining execution costs indicating computational costs for executing the task on the trained machine-learning model **702** or the third-party machine-learning model(s) **706**. For example, the intelligent selection and execution platform **102** determines a compute metric indicating the amount of computational power required for executing the task on the trained machine-learning model **702** or the third-party machine-learning model(s) **706**. As another example, the intelligent selection and execution platform **102** can determine an execution cost metric by determining an amount of CPU or GPU required by the trained machine-learning model **702** or the third-party machine-learning model(s) **706** to execute the task.

(126) Further, as shown in FIG. 7, in one or more embodiments, the intelligent selection and execution platform **102** can determine task routing metrics **708** by determining model fit metrics **716**. In particular, the intelligent selection and execution platform **102** determines model fit metrics **716** by determining model alignment for executing the task on the trained machine-learning model **702** and/or the third-party machine-learning model(s) **706**. For example, the intelligent selection and execution platform **102** determines if the trained machine-learning model **702** or the third-party machine-learning model(s) **706** comprise an amount of computational power that is more than necessary for executing the task (e.g., using a larger machine-learning model for a small task) or if they can execute the task but comprises a smaller amount of computational power (e.g., using a smaller machine-learning model for a large task).

(127) In addition, as illustrated in FIG. 7, in some embodiments, the intelligent selection and execution platform **102** can determine task routing metrics **708** by determining user feedback metrics **718**. In particular, the intelligent selection and execution platform **102** determines user feedback metrics **718** by accessing user feedback data comprising stored quality indicators specific to each model of the plurality of machine-learning models. As previously mentioned, the intelligent selection and execution platform **102** can receive user feedback data indicating user satisfaction with the performance of a machine-learning model or execution of the task. The intelligent selection and execution platform **102** can access the user feedback data and determine user feedback metrics **718** that indicate an estimated user satisfaction of a task executed on the trained machine-learning model **702** or the third-party machine-learning model(s) **706** (e.g., indicating whether a user will be satisfied with a task).

(128) As also illustrated in FIG. 7, in some embodiments, the intelligent selection and execution platform **102** determines task routing metrics **708** by determining model states **720**. Specifically, the intelligent selection and execution platform **102** determines model states **720**, which indicate whether the trained machine-learning model **702** or the third-party large-language model(s) **706** are available to execute a task. For example, model states **720** can comprise inferences or inference times of the trained machine-learning model **702** or the third-party large-language model(s) **706** indicating whether or not the trained machine-learning model **702** or the third-party large-language model(s) **706** are able to execute tasks. To illustrate, an inference time may indicate that the trained machine-learning model **702** or the third-party large-language model(s) **706** are executing tasks, and, as a result, response times from each respective model may not meet a response time threshold.

(129) Moreover, as illustrated in FIG. 7, the intelligent selection and execution platform **102** can determine task routing metrics **708** by determining model capabilities **722**. In particular, the intelligent selection and execution platform **102** determines model capabilities **722** by determining that the trained machine-learning model **702** or the third-party large-language model(s) **706** are able to perform certain actions. For example, model capabilities **722** indicates that, if provided a request to execute a certain task, the trained machine-learning model **702** or the third-party large-language model(s) **706** are capable of executing the task. To illustrate, model capabilities **722** can indicate that the trained machine-learning model **702** or the third-party large-language model(s) **706** are capable of receiving a certain kind of input, generating a certain kind of output, executing a certain task, or executing tasks of a certain size (e.g., large tasks). In some cases, the intelligent selection and execution platform **102** identifies that a task is a hot path task that requires continuous execution and will utilize a model capability metric to align the task with the hot path task.

(130) Also, as shown in FIG. 7, in one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics **708** by determining model specialties **724**. Specifically, the intelligent selection and execution platform **102** can determine model specialties **724** that indicate that the trained machine-learning model **702** or the third-party large-language model(s) **706** are adept at performing a certain task or generating a certain output. For example,

model specialties **724** indicate that the trained machine-learning model **702** or the third-party large-language model(s) **706** perform a certain task or generate a certain output a measurable amount better than other machine-learning models. To illustrate, the intelligent selection and execution platform **102** identifies model specialties **724** based on known specialties (e.g., a machine-learning model was generated or trained to perform a certain task). As another illustration, the intelligent selection and execution platform **102** can generate quality metrics that indicate model specialties **724** for trained machine-learning model **702** or the third-party large-language model(s) **706** based on the quality metrics. Moreover, in one or more embodiments, the intelligent selection and execution platform **102** learns model specialties **724** based on a software domain analysis of the plurality of machine-learning models.

(131) As previously mentioned, in some embodiments, the intelligent selection and execution platform **102** utilizes a model selection machine-learning model. In particular, the intelligent selection and execution platform **102** utilizes the model selection machine-learning model to select, from a plurality of machine-learning models, a designated machine-learning model for executing a task. FIG. **8** illustrates a schematic diagram of an intelligent selection and execution platform utilizing a model selection machine-learning model to select a designated machine-learning model for executing a task in accordance with one or more embodiments.

(132) As also previously mentioned, the intelligent selection and execution platform **102** selects a designated machine-learning model based on workload features defining characteristics of a task and task routing metrics for a plurality of machine-learning models. As shown in FIG. **8**, model selection machine-learning model **808** utilizes task routing metrics **804** and workload features **806** to select designated machine-learning model **810**. Additional details regarding the intelligent selection and execution platform **102** extracting workload features are provided in relation to FIG. **6** above. Additional details regarding the intelligent selection and execution platform **102** determining task routing metrics for a plurality of machine-learning models are provided in relation to FIG. **7** above.

(133) As illustrated in FIG. **8**, in one or more embodiments, model selection machine-learning model **808** receives or performs software domain analysis **802**. In particular, the intelligent selection and execution platform **102** performs a software domain analysis by analyzing the plurality of machine-learning models to find commonalities and variables between the machine-learning models. For example, as part of a software domain analysis, the intelligent selection and execution platform **102** can identify that a machine-learning model has a specialty in a certain aspect of executing the task and utilize the software domain analysis when selecting a designated machine-learning model for the task. In some cases, intelligent selection and execution platform **102** performs a software domain analysis **802** and stores data from the software domain analysis **802** for recall at another time. In other cases, the intelligent selection and execution platform **102** trains the model selection machine-learning model **808** based on the results of the software domain analysis **802**.

(134) As previously mentioned, the intelligent selection and execution platform **102** selects a designated machine-learning model based on determining an optimal machine-learning model for executing a task. In particular, the model selection machine-learning model **808** is trained, tuned, or optimized to weigh workload features and task routing metrics to determine an optimal machine-learning model. In some cases, the model selection machine-learning model **808** is a multi-armed bandit model that is trained or optimized to select a designated machine-learning model from a plurality of machine-learning models. For example, the intelligent selection and execution platform **102** trains or optimizes a multi-armed bandit model to select a designated machine-learning model based on workload data and task routing metrics of the plurality of machine-learning models at the time of selection. In addition, as described further below in relation to FIGS. **10A-10B**, a multi-armed bandit model continues to improve over time as it receives additional feedback (e.g., user feedback data).

(135) As mentioned, the model selection machine-learning model **808** determines an optimal machine-learning model for executing a task. Specifically, the model selection machine-learning model **808** generates an optimization metric for each machine-learning model of a plurality of machine-learning models to determine an optimal machine-learning model. For example, the model selection machine-learning model **808** can analyze workload features **806**, task routing metrics **804**, and software domain analysis **802** to generate an optimization metric for each machine-learning model, and the intelligent selection and execution platform **102** determines an optimal machine-learning model based on the optimization metric.

(136) In some cases, the intelligent selection and execution platform **102** selects an optimal model based on the optimization metric satisfying an optimal model threshold. Specifically, the intelligent selection and execution platform **102** determines that a machine-learning model will execute the task at a satisfactory level based on an optimization metric for the machine-learning model satisfying the optimal model threshold. For example, when an optimization metric satisfies the optimal model threshold, a machine-learning model is likely to execute the task in a threshold amount of time or satisfy a quality metric. As another example, an optimization metric may satisfy the optimal model threshold if the machine-learning model comprises a model capability or model specialty associated with the task (e.g., the capability or specialty aligns with the requested task).

(137) In addition, in one or more embodiments, the intelligent selection and execution platform **102** selects a designated machine-learning model based on a historical quality metric. In particular, the intelligent selection and execution platform **102** utilizes user feedback to determine a historical quality metric that indicates an estimated quality of output from the machine-learning model. The intelligent selection and execution platform **102** then utilizes the historical quality metric to select the designated machine-learning model.

(138) As also shown in FIG. **8**, the model selection machine-learning model **808** selects a designated machine-learning model **810** for executing task **812**. In particular, the model selection machine-learning model **808** selects the designated machine-learning model **810**, and the intelligent selection and execution platform **102** provides a data to the designated machine-learning model **810** to execute the task **812**. Example tasks include generating summaries, generating documents, extracting knowledge and/or relations in text, generating code, analyzing sentiments, or parsing documents. In some cases, the intelligent selection and execution platform **102** utilizes content items stored in the content management system **108** to execute tasks, such as generating summaries for content items or generating documents based on content items.

(139) In one or more embodiments, the intelligent selection and execution platform **102** selects a large language model as designated machine-learning model **810**. In particular, the intelligent selection and execution platform **102** identifies that a large language model is necessary to execute task **812** and provides a prompt to the large language model to execute task **812**. In some cases, the intelligent selection and execution platform **102** can identify that workload data requests execution of a task using a large language model and selects a designated large language model from a plurality of large language models.

(140) In one or more embodiments, the intelligent selection and execution platform **102** can also generate a machine-learning model. In particular, the intelligent selection and execution platform **102** can determine that generating a machine-learning model for executing the task requires less computational and/or hardware cost than selecting an existing model and generating a model that is specific to the task and generate a machine-learning model for executing the task.

(141) As also illustrated in FIG. **8**, the intelligent selection and execution platform **102** also selects a designated data storage **814**. In particular, the intelligent selection and execution platform **102** utilizes the model selection machine-learning model **808** to determine designated data storage **814** to use in conjunction with executing task **812**. In particular, the intelligent selection and execution platform **102** determines the financial cost of utilizing databases and other storage systems within the intelligent selection and execution platform **102** or sending the data to a third-party database or

storage system. For example, the intelligent selection and execution platform **102** determines a currency amount required to store data at third-party databases and databases within the intelligent selection and execution platform **102**. To illustrate, the intelligent selection and execution platform **102** can determine that sending data to a third-party database or storage system has a higher financial cost (e.g., higher currency) than storing the data within the intelligent selection and execution platform **102**.

(142) In addition, the intelligent selection and execution platform **102** can determine an execution cost when determining designated data storage **814**. In particular, the intelligent selection and execution platform **102** determines an execution cost of utilizing databases and other storage systems within the intelligent selection and execution platform **102** or sending the data to a third-party database or storage system. For example, the intelligent selection and execution platform **102** can determine computational costs, such as processing or memory, the executional cost to store data at third-party databases, and the executional cost to store data at a database within the intelligent selection and execution platform **102**. To illustrate, the intelligent selection and execution platform **102** can determine that an amount of time required to send data to a third-party database or storage system is higher than a threshold (e.g., it is slower than storing within the intelligent selection and execution platform **102** or the content management system).

(143) As illustrated in FIG. **8**, the intelligent selection and execution platform **102** receives user feedback data **816** based on the execution of task **812**. In particular, the intelligent selection and execution platform **102** receives user feedback data that indicates a quality measure of execution of the task. Further, as indicated in FIG. **8**, the intelligent selection and execution platform **102** can update parameters of the model selection machine-learning model **808** based on user feedback data **816**. Additional details regarding the intelligent selection and execution platform **102** receiving user feedback data and training or updating the parameters of the model selection machine-learning model **808** are provided in relation to FIGS. **10A-10B** below.

(144) As previously mentioned, the intelligent selection and execution platform **102** flexibly adds machine-learning models to a plurality of machine-learning models. In particular, the intelligent selection and execution platform **102** updates a model selection machine-learning model to add additional machine-learning models to a plurality of machine-learning models. FIG. **9** illustrates a schematic diagram of an intelligent selection and execution platform **102**, adding an additional machine-learning model to a plurality of machine-learning models in accordance with one or more embodiments.

(145) In one or more embodiments, the intelligent selection and execution platform **102** adds additional machine-learning model **902** by providing or accessing task routing metrics for the additional machine-learning model **902**. As illustrated in FIG. **9**, model selection machine-learning model **912** receives offline task routing metrics **904** for additional machine-learning model **902**. In particular, offline task routing metrics **904** comprise scores or other measured data that indicate the quality of additional machine-learning model **902** and that intelligent selection and execution platform **102** can utilize to evaluate additional machine-learning model **902**. For example, offline task routing metrics **904** comprise task routing metrics that are measured, scored, or annotated with ground truth quality labels. In some cases, the ground truth quality labels are provided or confirmed by human observers to indicate measures of quality for the task routing metrics. The intelligent selection and execution platform **102** provides the offline task routing metrics **904** with the model selection machine-learning model **912** to determine that the additional machine-learning model **902** produces output above a threshold quality.

(146) In addition, the intelligent selection and execution platform **102** receives online task routing metrics **906**. In particular, online task routing metrics **906** comprise task routing metrics utilized during (or as a part of) model selection machine-learning model selecting a designated machine-learning model. For example, online quality metrics comprise user feedback data received upon execution of a task and as described in further detail in FIG. **10A-10B** below. In one or more

additional embodiments, online task routing metrics comprise quality indicators specific to the task. For example, model selection machine-learning model **912** accesses the online quality indicators when selecting a designated machine-learning model for a task.

(147) As also illustrated, the intelligent selection and execution platform **102** can also utilize properties **908** for the additional machine-learning model. In particular, the intelligent selection and execution platform **102** accesses properties **908** that indicate sizes, data types, or additional information for additional machine-learning model **902**. For example, properties **908** can indicate the size of additional machine-learning model **902**. As another example, properties **908** can include the data type in which additional machine-learning model **902** was trained (e.g., Brain Float, 8-bit). Moreover, as an example, properties **908** can indicate the benchmarks and evaluations for additional machine-learning model **902**. Further, properties **908** can indicate the layers inside the model and which generations of GPU can support those layers. Also, properties **908** can denote a context window size for additional machine-learning model **902**.

(148) As further illustrated, the intelligent selection and execution platform **102** can also utilize characteristics **910** for the additional machine-learning model **902**. In particular, characteristics **910** indicate model capabilities or model specialties of additional machine-learning model **902**. For example, characteristics **910** can indicate that additional machine-learning model **902** is capable of executing a certain type of task, a task of a certain size, or on a certain type of hardware. Moreover, characteristics **910** can indicate that additional machine-learning model **902** comprises a certain specialty, was designed to execute a certain task, or was designed to access a certain hardware environment.

(149) As illustrated, the intelligent selection and execution platform **102** adds the additional language model to the plurality of machine-learning models **914**. In particular, once the intelligent selection and execution platform **102** identifies that the additional machine-learning model performs within a threshold effectiveness and can access task routing metrics for the additional machine-learning model, the intelligent selection and execution platform **102** adds the additional machine-learning model to the plurality of machine-learning models **914** to establish an updated plurality of machine-learning models. Moreover, because the model selection machine-learning model **912** is continuously being updated (e.g., through user feedback data), the model selection machine-learning model **912** adds the machine-learning model without time-intensive and computationally expensive training. Indeed, the model selection machine-learning model **912** will continue to improve and adapt to the additional machine-learning model **902** over time and based on additional user feedback data.

(150) In addition, though not illustrated in FIG. 9, the intelligent selection and execution platform **102** can also add additional hardware environments to a plurality of hardware environments using a hardware allocating machine-learning model. Similar to adding an additional machine-learning model **902**, the intelligent selection and execution platform **102** receives offline task routing metrics, online task routing metrics, properties, and characteristics for an additional hardware environment. The intelligent selection and execution platform **102** can analyze additional hardware environments with the offline task routing metrics, online task routing metrics, properties, and characteristics and, upon determining that the additional hardware environment can execute the task within a threshold quality, add the additional hardware environment to the plurality of hardware environments. Similar to the model selection machine-learning model, the hardware allocating machine-learning model will continue to improve selections of the additional hardware environment based on receiving continuous feedback regarding executing the task using the hardware environment.

(151) As previously mentioned, the intelligent selection and execution platform **102** receives user feedback data. In particular, the intelligent selection and execution platform **102** utilizes user feedback data to update parameters of a model-selection machine-learning model. FIGS. 10A-10B illustrate an example diagram of an intelligent selection and execution platform receiving user

feedback data and utilizing the user feedback data to update parameters of a model selection machine-learning model in accordance with one or more embodiments.

(152) As illustrated in FIG. 10A, after executing task **1002**, the intelligent selection and execution platform **102** receives user feedback data **1004**. Specifically, the intelligent selection and execution platform **102** receives user feedback data **1004** by receiving information, opinions, or responses that indicate user satisfaction with the task or how the machine-learning model executed the task. For example, user feedback data can be received continuously or in real-time as tasks are executed by a designated machine-learning model selected by the intelligent selection and execution platform **102** (or the model selection machine-learning model). In some cases, the intelligent selection and execution platform **102** can determine user feedback data **1004** is positive user feedback data indicating user satisfaction with the output of the machine-learning model. In other cases, the intelligent selection and execution platform **102** can determine that user feedback data **1004** is negative user feedback data indicating a user dissatisfaction with the output of the machine-learning model.

(153) As illustrated, in one or more embodiments, user feedback data **1004** can comprise explicit feedback data **1006**. Specifically, the intelligent selection and execution platform **102** receives explicit feedback by receiving user input indicating user satisfaction with the task or with how the machine-learning model executed task **1002**. For example, after the designated (or fallback) machine-learning model executes task **1002**, intelligent selection and execution platform **102** provides an option to provide user feedback data on the client device that provided the request to execute the task (e.g., through workload data). In some cases, the intelligent selection and execution platform **102** receives positive explicit feedback or negative explicit feedback. Additional details regarding the intelligent selection and execution platform **102** providing an option to provide user feedback data are provided in relation to FIG. 10B below.

(154) As further illustrated, in some embodiments, user feedback data **1004** can comprise implicit feedback data **1008**. In particular, implicit feedback refers to user feedback data that is suggested, implied, or understood without being explicitly provided by the client device. For example, implicit feedback measures or infers user satisfaction with a task by monitoring user interactions or client device behavior after receiving the output from the machine-learning model. Implicit feedback data **1008** can refer to an action performed by the client device indicating user satisfaction with the output. To illustrate, if a designated machine-learning model provides a list of results to the client device and suggests that the client device copy the list of result and the client device copies the list of results, the intelligent selection and execution platform **102** will identify positive implicit feedback indicating user satisfaction with the output of the designated machine-learning model.

(155) Conversely, implicit feedback data **1008** can indicate a user dissatisfaction with the output based on identifying that the client device did not take an action suggested in the output of the machine-learning model. To illustrate, if a designated machine-learning model provides a list of results to the client device and suggests that the client device copy the list of results and the client device chooses not to copy the list of results, the intelligent selection and execution platform **102** will identify negative implicit feedback indicating a user dissatisfaction with the output of the designated machine-learning model. As another illustration, if a designated machine-learning model provides an output and the client device provides the same prompt to the intelligent selection and execution platform **102** again, then the intelligent selection and execution platform **102** will identify negative implicit feedback indicating a user dissatisfaction with the output of the designated machine-learning model.

(156) As illustrated in FIG. 10A, the intelligent selection and execution platform **102** provides user feedback data **1004** to model selection machine-learning model **1010**. In particular, the intelligent selection and execution platform **102** continuously (or in real-time) provides user feedback data **1004** to model selection machine-learning model **1010** in order to continuously update parameters of the model selection machine-learning model. For example, the model selection machine-learning

model **1010** can identify that output from a machine-learning model elicits positive user feedback data or negative user feedback data from the output of the designated machine-learning model and updates parameters of the model selection machine-learning model **1010**. To illustrate, if the model selection machine-learning model **1010** selected a machine-learning model with a lower financial cost metric and a smaller model fit metric (e.g., a smaller, less powerful machine-learning model compared to other machine-learning models) and received user negative user feedback data, the model selection machine-learning model may select a different machine-learning model when presented with a prompt to execute a similar task.

(157) In one or more embodiments, the intelligent selection and execution platform **102** can generate quality metrics based in part on the user feedback data **1004**. Specifically, the intelligent selection and execution platform **102** generates quality metrics that can be used to compare the quality of machine-learning models. For example, the intelligent selection and execution platform **102** can utilize the quality metrics as task routing metrics for selecting a designated machine-learning model or adding an additional machine-learning model to the plurality of machine-learning models. To illustrate, if user feedback data **1004** indicates user satisfaction with an output type from a machine-learning model, the model selection machine-learning model **1010** may add the output type as a model specialty of the machine-learning model and utilize the machine-learning model when receiving a task that requests the output type.

(158) As previously mentioned, the intelligent selection and execution platform **102** can receive explicit user feedback data by receiving user feedback data after a designated machine-learning model executes a task. FIG. **10B** illustrates an example graphical user interface **1012**. Specifically, the graphical user interface **1012** is rendered on a client device by application of the content management system **108** or by a third-party application connected to the content management system **108**. For example, graphical user interface **1012** renders or displays the output of the designated machine-learning model on the client device that provided workload data requesting execution of the task.

(159) As illustrated, in some embodiments, the intelligent selection and execution platform **102** can provide a user feedback data request **1014** in the graphical user interface **1012**. In particular, the intelligent selection and execution platform **102** can provide the user feedback data request after a designated machine-learning model executes the task. For example, the intelligent selection and execution platform **102** can detect whether or not the client device interacted with the output of the designated machine-learning model and provide the user feedback data request **1014**.

(160) As previously mentioned, the intelligent selection and execution platform **102** receives explicit feedback data **1006**. In particular, the intelligent selection and execution platform **102** receives explicit user feedback data by receiving user input in the user feedback data request **1014**. For example, as shown, a client device may provide user feedback data by selecting an option in the user feedback data request **1014** or by providing text input in the user feedback data request **1014**.

(161) FIGS. **1-10**, the corresponding text, and the examples provide a number of different methods, systems, devices, and non-transitory computer-readable media of the intelligent selection and execution platform **102**. In addition to the foregoing, one or more embodiments can also be described in terms of flowcharts comprising acts for accomplishing a particular result, as shown in FIG. **11**. FIG. **11** may be performed with more or fewer acts. Further, the acts may be performed in differing orders. Additionally, the acts described herein may be repeated or performed in parallel with one another or parallel with different instances of the same or similar acts.

(162) As mentioned, FIG. **11** illustrates a flowchart of a series of acts **1100** for selecting a designated machine-learning model in accordance with one or more embodiments. While FIG. **11** illustrates acts according to one embodiment, alternative embodiments may omit, add to, reorder, and/or modify any of the acts shown in FIG. **11**. The acts of FIG. **11** can be performed as part of a method. Alternatively, a non-transitory computer-readable medium can comprise instructions that,

when executed by one or more processors, cause a computing device to perform the acts of FIG.

11. In some embodiments, a system can perform the acts of FIG. **11.**

(163) As shown in FIG. **11**, the series of acts **1100** includes an act **1102** of receiving workload data requesting execution of a task using a machine-learning model, an act **1104** of extracting workload features defining characteristics of the task, an act **1106** of determining task routing metrics for a plurality of machine-learning models hosted in respective network environments, and an act **1108** of selecting a designated machine-learning model for executing the task based on the workload features and the task routing metrics.

(164) In particular, the act **1102** can include receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model, the act **1104** can include extracting, from the workload data, workload features defining characteristics of the task, the act **1106** can include determining task routing metrics for a plurality of machine-learning models hosted in respective network environments, and the act **1108** can include selecting, from the plurality of machine-learning models, a designated machine-learning model for executing the task based on the workload features and the task routing metrics.

(165) For example, in one or more embodiments, the series of acts **1100** includes wherein determining task routing metrics for the plurality of machine-learning models further comprises determining a model state for each machine-learning model of the plurality of machine-learning models and wherein selecting the designated machine-learning model for executing the task is based in part on the model state of the designated machine-learning model. In addition, in one or more embodiments, the series of acts **1100** includes wherein determining task routing metrics for the plurality of machine-learning models further comprises determining a financial cost metric, an execution time metric, an execution cost metric, or a model fit metric for executing the task on each machine-learning models of the plurality of machine-learning models, and wherein selecting the designated machine-learning model is based on two or more of the financial cost metric, the execution time metric, the execution cost metric, and the model fit metric.

(166) In addition, in one or more embodiments, the series of acts **1100** includes wherein selecting the designated machine-learning model further comprises analyzing the workload data and the task routing metrics for the plurality of machine-learning models to determine an optimal machine-learning model for executing the task from the plurality of machine-learning models, and selecting the designated machine-learning model for executing the task based on determining that the designated machine-learning model is the optimal machine-learning model for executing the task.

(167) Moreover, in one or more embodiments, the series of acts **1100** includes wherein selecting the designated machine-learning model further comprises analyzing the workload data and the task routing metrics for the plurality of machine-learning models to determine an optimal machine-learning model for executing the task from the plurality of machine-learning models and selecting the designated machine-learning model for executing the task based on determining that the designated machine-learning model is the optimal machine-learning model for executing the task. In one or more embodiments, the series of acts **1100** includes wherein determining that the designated machine-learning model is the optimal machine-learning model further comprises generating, utilizing a model selection machine-learning model, an optimization metric and determining that the designated machine-learning model is the optimal machine-learning model for executing the task based on the optimization metric.

(168) Also, in one or more embodiments, the series of acts **1100** includes wherein selecting the designated machine-learning model further comprises performing a software domain analysis of each machine-learning model of the plurality of machine-learning models for executing the task and selecting the designated machine-learning model based on the software domain analysis.

Moreover, in one or more embodiments, the series of acts **1100** includes wherein extracting workload features defining characteristics of the task further comprises determining an estimated processing requirement and an estimated storage requirement for executing the task.

(169) Further, in one or more embodiments, the series of acts **1100** includes adding an additional machine-learning model to the plurality of machine-learning models to establish an updated plurality of machine-learning models and selecting the designated machine-learning model from the updated plurality of machine-learning models. In addition, in one or more embodiments, the series of acts **1100** includes accessing user feedback metrics about executing tasks using one or more machine-learning models of the plurality of machine-learning models, generating a historical quality metric based on the user feedback metrics, and selecting the designated machine-learning model for the task based on the historical quality metric.

(170) Also, in one or more embodiments, the series of acts **1100** includes wherein determining task routing metrics for the plurality of machine-learning models further comprises identifying a capability or a specialty associated with one or more machine-learning models of the plurality of machine-learning models and wherein selecting the designated machine-learning model for executing the task is based on alignment of the capability or the specialty of the designated machine-learning model with the task.

(171) In addition, in one or more embodiments, the series of acts **1100** includes receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model, extracting, from the workload data, workload features defining characteristics of the task, determining task routing metrics for a plurality of machine-learning models comprising one or more trained machine-learning models and one or more third-party machine-learning models hosted in respective network environments, and selecting, from the plurality of machine-learning models, a designated machine-learning model for executing the task based on the workload features and the task routing metrics.

(172) Moreover, in one or more embodiments, the series of acts **1100** includes identifying, based on a model state for the designated machine-learning model, that the designated machine-learning model is unavailable and selecting a fallback machine-learning model from the plurality of machine-learning models for executing the task. In addition, in one or more embodiments, the series of acts **1100** includes selecting the designated machine-learning model by determining, based on the task routing metrics, a financial cost metric for executing the task on each machine-learning model of the plurality of machine-learning models and selecting the designated machine-learning model based in part on the financial cost metric. In addition, the series of acts **1100** includes determining, based on the workload data and the task routing metrics, that a trained machine-learning model of the one or more trained machine-learning models is an optimal model for executing the task, identifying that the trained machine-learning model is unavailable, and selecting the designated machine-learning model from the one or more third-party machine-learning models based on identifying that the trained machine-learning model is unavailable.

(173) Further, in one or more embodiments, the series of acts **1100** includes selecting the designated machine-learning model by identifying, based on the workload data, that executing the task requires a machine-learning model comprising a capability or a specialty, identifying that a third-party machine-learning model of the one or more third-party machine-learning models comprises the capability or the specialty, and selecting the third-party machine-learning model as the designated machine-learning model for executing the task based on alignment of the capability or the specialty of the third-party machine-learning model with the task.

(174) In addition, in one or more embodiments, the series of acts **1100** includes receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model, extracting, from the workload data, workload features defining characteristics of the task, determining task routing metrics for a plurality of machine-learning models hosted in respective network environments, and selecting, utilizing a model selection machine-learning model, a designated machine-learning model from the plurality of machine-learning models for executing the task based on the workload features and the task routing metrics. Moreover, in one or more embodiments, the series of acts **1100** includes utilizing

the model selection machine-learning model to compare one or more machine-learning models of the plurality of machine-learning models based on the workload features and the task routing metrics, where the plurality of machine-learning models comprises one or more trained machine-learning models and one or more third-party trained machine-learning models and selecting the designated machine-learning model from the plurality of machine-learning models based on an output of the model selection machine-learning model.

(175) Also, in one or more embodiments, the series of acts **1100** includes receiving user feedback data indicating a user satisfaction with performance of the designated machine-learning model and updating parameters of the model selection machine-learning model based on the user feedback data. In addition, in one or more embodiments, the series of acts **1100** includes receiving additional task routing metrics for an additional machine-learning model and updating parameters of the model selection machine-learning model based on the additional task routing metrics of the additional machine-learning model. Moreover, in one or more embodiments, the series of acts **1100** includes generating, utilizing the model selection machine-learning model, an optimization metric for each machine-learning model of the plurality of machine-learning models and determining that the designated machine-learning model is an optimal machine-learning model for executing the task based on the optimization metric.

(176) As previously mentioned, the intelligent selection and execution platform **102** can select a hardware environment for executing a task. In particular, the intelligent selection and execution platform **102** can extract workload features of a task, determine task routing metrics for various hardware environments, and use the workload features and task routing metrics to select a hardware environment for executing the task. FIG. **12** illustrates a schematic diagram of an intelligent selection and execution platform selecting a hardware environment for executing a task in accordance with one or more embodiments.

(177) As illustrated, the intelligent selection and execution platform **102** performs an act **1202** and receives workload data requesting the execution of a task. In particular, the intelligent selection and execution platform **102** receives workload data from a device connected to the content management system that requests the execution of a task. For example, workload data comprises the necessary data and/or information for executing the task, such as a request to execute the task using a machine-learning model. In one or more embodiments, the intelligent selection and execution platform **102** receives workload data through an API layer, as described in relation to FIG. **4** above.

(178) In one or more embodiments, the intelligent selection and execution platform **102** workload data comprises parameters for executing a task. Specifically, the workload data comprises parameters comprising requests or specifications for executing a task. For example, workload data can comprise a parameter that instructs the intelligent selection and execution platform **102** to utilize a certain machine-learning model for executing a task or to execute the task using a certain hardware environment. As another example, workload data can comprise a parameter that instructs the intelligent selection and execution platform **102** to not utilize a certain machine-learning model or hardware environment to execute a task. Additional detail regarding the intelligent selection and execution platform **102** receiving workload data and parameters for executing a task is provided in relation to FIG. **6** above

(179) As also shown in FIG. **12**, the intelligent selection and execution platform **102** performs an act **1204** and extracts, from the workload data, workload features defining characteristics of the task. In particular, the intelligent selection and execution platform **102** extracts workload features by determining an amount of processing and storage needed for executing the task. For example, the intelligent selection and execution platform **102** can compute the computational cost for executing a task using a specific hardware environment. In some cases, the intelligent selection and execution platform **102** determines processing and storage requirements for executing a task on a hardware environment of the content management system and third-party hardware environments. Additional details regarding the intelligent selection and execution platform **102** extracting

workload features from workload data are provided in relation to FIG. 6 above.

(180) As further shown in FIG. 12, the intelligent selection and execution platform **102** performs an act **1206** and determines task routing metrics for a plurality of hardware environments hosted in respective network environments. In particular, the intelligent selection and execution platform **102** determines task routing metrics for a hardware environment within the intelligent selection and execution platform **102** and third-party hardware environments located in their respective network environments. For example, the intelligent selection and execution platform **102** can identify task routing metrics that indicate financial costs, execution times, and execution time for executing the task using the hardware environment of the intelligent selection and execution platform **102** or third-party hardware environments. In addition, the intelligent selection and execution platform **102** determines task routing metrics by determining a hardware state indicating whether the hardware environment is available to execute the task. Further, intelligent selection and execution platform **102** determines or identifies hardware environment capabilities or specialties. Additional details regarding the intelligent selection and execution platform **102** determining task routing metrics for a plurality of hardware environments are provided in relation to FIG. 13 below.

(181) As also shown in FIG. 12, the intelligent selection and execution platform **102** performs an act **1208** and selects a hardware environment for executing a task. In particular, the intelligent selection and execution platform **102** utilizes workload features and task routing metrics to select a designated hardware environment for executing the task. For example, the intelligent selection and execution platform **102** selects a hardware environment for executing the task by using a trained model, such as a neural network or a multi-armed bandit model, to determine the hardware environment for executing the task. In one or more embodiments, the intelligent selection and execution platform **102** utilizes a hardware allocating machine-learning model to select a designated hardware environment for executing the task. In some cases, the hardware allocating machine-learning model is trained, tuned, or optimized to select a designated hardware environment for executing a task.

(182) As mentioned, in one or more embodiments, the intelligent selection and execution platform **102** selects a designated hardware environment for executing a task based on workload features and task routing metrics. Specifically, the intelligent selection and execution platform **102** utilizes a hardware allocating machine-learning model to analyze the workload features and task routing metrics and select a designated hardware environment for executing the task. For example, the hardware allocating machine-learning model generates an optimization metric based on the workload features and the task routing metrics and selects the designated hardware environment based on the optimization metric. Additional details regarding the intelligent selection and execution platform **102** utilizing a hardware allocating machine-learning model to select a designated hardware environment for executing a task are discussed further in FIG. 14 below.

(183) As previously mentioned, in one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics for a plurality of hardware environments. In particular, the intelligent selection and execution platform **102** determines task routing metrics for a hardware environment of the intelligent selection and execution platform **102** and one or more third-party hardware environments. FIG. 13 illustrates a schematic diagram of an intelligent selection and execution platform determining task routing metrics for a plurality of hardware environments in accordance with one or more embodiments.

(184) As illustrated in FIG. 13, the intelligent selection and execution platform **102** determines task routing metrics **1308**. Specifically, the intelligent selection and execution platform **102** determines task routing metrics **1308** that indicate costs, timeframes, capabilities, fit, or feedback for each hardware environment of a plurality of hardware environments. For example, the intelligent selection and execution platform **102** determines task routing metrics for a hardware environment **1302** of the intelligent selection and execution platform **102** (e.g., local to the intelligent selection and execution platform **102** or the content management system **108**) and for third-party hardware

environment(s) **1306** located on third-party server(s) **1304**. As indicated, the intelligent selection and execution platform **102** can determine task routing metrics for multiple third-party hardware environments in their respective network environments (e.g., on their own servers and/or with their own network).

(185) In one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics **1308** for executing the task with different machine-learning models using hardware environment **1302** and/or third-party hardware environment(s) **1306**. Specifically, the intelligent selection and execution platform **102** can determine task routing metrics for different combinations of machine-learning models and hardware environment **1302** and/or third-party hardware environment(s) **1306**. To illustrate, the intelligent selection and execution platform **102** determines task routing metrics **1308** for executing a task using the same machine-learning model on the hardware environment **1302** and third-party hardware environment(s) **1306** or by using different machine-learning models on the hardware environment **1302** and/or the third-party hardware environment(s) **1306**. Indeed, the intelligent selection and execution platform **102** determines task routing metrics **1308** for executing tasks using various machine-learning models on various hardware environments to compare availabilities, capabilities, costs, and fits for the task and to select a designated machine-learning model and/or hardware environment for executing the task.

(186) As shown, in one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics **1308** by determining financial cost metrics **1310**. In particular, the intelligent selection and execution platform **102** can determine financial cost metrics **1310** by determining financial costs for executing the task using each of hardware environment **1302** and the third-party hardware environment(s) **1306**. For example, the intelligent selection and execution platform **102** can determine a currency cost for executing a task on each of the hardware environment **1302** and the third-party hardware environment(s) **1306**. Moreover, in one or more embodiments, the intelligent selection and execution platform **102** can determine financial cost metrics **1310** based on a use time period. For example, financial cost metrics **1310** may vary based on bandwidth availability at certain points throughout a use time period (e.g., different times of day). To illustrate, executing a task on a third-party hardware environment(s) may be higher at a high-use (or peak) time period than executing the task during a minimum-use time period.

(187) As also shown in FIG. **13**, the intelligent selection and execution platform **102** determines task routing metrics **1308** by determining execution time metrics **1312**. In particular, the intelligent selection and execution platform **102** determines execution time metrics **1312** by determining an execution time for executing a task on hardware environment **1302** and/or third-party hardware environment(s) **1306**. For example, the intelligent selection and execution platform **102** can determine that hardware environment **1302** will execute a task in a first time frame and third-party hardware environment(s) **1306** will execute a task in a second time frame. To illustrate, the intelligent selection and execution platform **102** can determine execution time metrics **1312** that indicate third-party hardware environment(s) **1306** will execute a task in less time than hardware environment **1302**.

(188) As further illustrated in FIG. **13**, the intelligent selection and execution platform **102** determines task routing metrics **1308** by determining execution cost metrics **1314**. In particular, the intelligent selection and execution platform **102** determines execution cost metrics **1314** by determining computational costs for executing the task on the hardware environment **1302** and the third-party hardware environment(s) **1306**. For example, the intelligent selection and execution platform **102** determines a compute metric indicating the computational power needed to execute the task on the hardware environment **1302** and/or the third-party hardware environment(s) **1306**. For example, the intelligent selection and execution platform **102** can determine an execution cost metric by determining an amount of processing power needed by a CPU or GPU to execute the task.

(189) In addition, as shown in FIG. 13, the intelligent selection and execution platform **102** determines task routing metrics **1308** by determining user feedback metrics **1316**. As previously mentioned, the intelligent selection and execution platform **102** can receive user feedback data indicating user satisfaction with how the machine-learning model executed the task, including user feedback data that indicates user satisfaction with how the hardware environment executed the task, such as when a user was dissatisfied with the processing time for executing the task. The intelligent selection and execution platform **102** can determine (or generate) user feedback metrics **1316** based on the user feedback data. For example, the intelligent selection and execution platform **102** determines (or generates) user feedback metrics **1316** based on tasks executed by the hardware environment **1302** and the third-party hardware environment(s) **1306**. Additional details regarding the intelligent selection and execution platform **102** receiving user feedback data are provided in relation to FIGS. **10A-10B** above.

(190) Moreover, as illustrated in FIG. 13, the intelligent selection and execution platform **102** determines task routing metrics **1308** by determining hardware environment state **1318**. Specifically, the intelligent selection and execution platform **102** determines a hardware environment state **1318** for the hardware environment **1302** and each of the third-party hardware environment(s) **1306** that indicates bandwidth availability for each respective hardware environment. For example, the intelligent selection and execution platform **102** determines a hardware environment state **1318** by determining processor availability, such as whether a hardware environment is being used by other processes. As another example, the intelligent selection and execution platform **102** can determine if the memory size available in the hardware environment is capable of executing the task (e.g., the memory amount available exceeds the memory required to execute the task).

(191) In addition, as illustrated in FIG. 13, the intelligent selection and execution platform **102** determines task routing metrics by determining hardware environment capabilities **1320**. In particular, the intelligent selection and execution platform **102** determines hardware environment capabilities **1320** by determining that hardware environment **1302** or third-party hardware environment(s) **1306** comprise the necessary hardware components to facilitate executing tasks. For example, hardware environment capabilities **1320** indicate CPU capabilities, GPU capabilities, memory capacities, or storage capacities of the hardware environment **1302** and/or the third-party hardware environment(s) **1306**. To illustrate, in some cases, in order to execute a task, a machine-learning model may require a certain CPU capacity, a certain GPU capacity, a certain memory capacity, or a certain storage capacity, and hardware environment capabilities **1320** indicate whether or not a given hardware environment aligns with a machine-learning model.

(192) In addition, as illustrated in FIG. 13, the intelligent selection and execution platform **102** determines task routing metrics **1308** by determining hardware environment specialties **1322**. In particular, the intelligent selection and execution platform **102** determines hardware environment specialties **1322** that indicate that hardware environment **1302** or third-party hardware environment(s) **1306** are specialized or optimized for a machine-learning model. For example, hardware environment specialties **1322** indicate that when a machine-learning model utilizes a certain hardware environment to execute the task, the hardware environment will facilitate the machine-learning model to execute the task in less time than another hardware environment.

(193) As previously mentioned, in some embodiments, the intelligent selection and execution platform **102** utilizes a hardware allocating machine-learning model. In particular, the intelligent selection and execution platform **102** utilizes the hardware allocating machine-learning model to select, from a plurality of hardware environments, a designated hardware environment to execute a task. FIG. 14 illustrates a schematic diagram of an intelligent selection and execution platform utilizing a model hardware allocating machine-learning model to select a designated hardware environment in accordance with one or more embodiments.

(194) As also previously mentioned, the intelligent selection and execution platform **102** selects a

designated hardware environment based on workload features defining characteristics of a task and task routing metrics for a plurality of hardware environments. As shown in FIG. 14, hardware allocating machine-learning model **1408** utilizes task routing metrics **1404** and workload features **1406** to select designated hardware environment **1410**. Additional details regarding the intelligent selection and execution platform **102** extracting workload features defining characteristics of a task are provided in relation to FIG. 6 above. Additional details regarding the intelligent selection and execution platform **102** determining task routing metrics for a plurality of hardware environments are provided in relation to FIG. 13 above.

(195) As also illustrated in FIG. 14, in one or more embodiments, hardware allocating machine-learning model **1408** receives hardware load analysis **1402**. In particular, the intelligent selection and execution platform **102** analyzes various hardware environments when receiving workload data in order to estimate (or determine) the current computing capacities of the plurality of hardware environments. In some instances, the intelligent selection and execution platform **102** uses sharding techniques to estimate hardware environment capacities. For example, the intelligent selection and execution platform **102** samples hardware environments to determine models that have the capacity to execute a task. To illustrate, the intelligent selection and execution platform **102** can sample third-party current capacities of two or more third-party hardware environments and compare the third-party current capacities to the current capacities of the hardware environment of the intelligent selection and execution platform **102**. The intelligent selection and execution platform **102** can then generate a likelihood that one of the third-party hardware environments or the hardware environment of the intelligent selection and execution platform **102** has sufficient capacity to execute a task. Indeed, by using sharding techniques, the intelligent selection and execution platform **102** utilizes less computing power to determine task routing metrics for each hardware environment of the plurality of hardware environments, which is computationally expensive.

(196) In addition to sharding techniques, in some instances, the intelligent selection and execution platform **102** performs hardware load analysis **1402** by performing probabilistic load balancing. Specifically, the intelligent selection and execution platform **102** performs probabilistic load balancing by using randomization and statistical techniques to determine how tasks are distributed across a plurality of hardware environments with the goal of balancing the tasks efficiently. For example, the intelligent selection and execution platform **102** utilizes probabilistic load balancing to prevent a single hardware environment from becoming a bottleneck. In addition, the intelligent selection and execution platform **102** can determine whether or not sharding techniques and/or probabilistic load balancing are feasible within particular hardware environments as part of hardware load analysis **1402**. Moreover, the intelligent selection and execution platform **102** can also perform a combination of sharding techniques and probabilistic load balancing as part of hardware load analysis **1402**.

(197) As previously mentioned, the intelligent selection and execution platform **102** selects a designated machine-learning model based on determining an optimal hardware environment for executing a task. In particular, the hardware allocating machine-learning model **1408** is trained, tuned, or optimized to weigh workload features and task routing metrics to determine an optimal hardware environment. In some cases, the hardware allocating machine-learning model **1408** is a multi-armed bandit model that is trained or optimized to select a designated hardware environment from a plurality of hardware environments. For example, the intelligent selection and execution platform **102** trains or optimizes a multi-armed bandit model to select a designated hardware environment for a machine-learning model to utilize to execute a task based on workload data and task routing metrics of the plurality of hardware environments at the time of selection. In addition, as described below, a multi-armed bandit model continues to improve over time as it receives additional feedback (e.g., user feedback data).

(198) As mentioned, the hardware allocating machine-learning model **1408** determines an optimal

machine-learning model for executing a task. Specifically, in one or more embodiments, the hardware allocating machine-learning model **1408** generates an optimization metric for each hardware environment of a plurality of hardware environments to determine an optimal hardware environment. For example, the hardware allocating machine-learning model can analyze workload features **1406**, task routing metrics **1404**, and hardware load analysis **1402** to generate an optimization metric for each machine-learning model, then the intelligent selection and execution platform **102** determines an optimal hardware environment based on the optimization metric.

(199) In some cases, the intelligent selection and execution platform **102** selects an optimal hardware environment based on the optimization metric satisfying an optimal hardware environment threshold. Specifically, based on the optimization metric satisfying the optimal hardware environment threshold, the intelligent selection and execution platform **102** determines that the hardware environments will execute the task at a satisfactory level. For example, when an optimization metric satisfies the optimal hardware environment threshold, a hardware environment is likely to execute the task in a threshold amount of time or satisfy a quality metric. As another example, a optimization metric may satisfy the optimal hardware threshold if the hardware environment comprises a hardware environment capability or hardware environment specialty for a task (e.g., the capability or specialty aligns with the requested task).

(200) As also shown in FIG. **14**, the hardware allocating machine-learning model **1408** selects a designated hardware environment for executing task **1412**. In particular, the hardware allocating machine-learning model selects designated hardware environment **1410**, then provides data or input to a machine-learning model (e.g., a prompt to a large language model), and utilizes designated hardware environment **1410** to provide the computing capacity to execute task **1412**.

(201) In addition, in one or more embodiments, the intelligent selection and execution platform **102** selects a designated hardware environment based on a historical quality metric. In particular, the intelligent selection and execution platform **102** utilizes user feedback to determine a historical quality metric that indicates an estimated quality of output from tasks that utilize the hardware environment. The intelligent selection and execution platform **102** then utilizes the historical quality metric to select the designated hardware environment.

(202) Further, as illustrated in FIG. **14**, the intelligent selection and execution platform **102** receives user feedback data **1414** based on the execution of task **1412**. In particular, the intelligent selection and execution platform **102** receives user feedback data that indicates a quality measure of execution of the task. Further, as indicated in FIG. **14**, the intelligent selection and execution platform **102** can update parameters of the hardware allocating machine-learning model **1408** based on user feedback data **1414**. Additional details regarding the intelligent selection and execution platform **102** receiving user feedback data and training or updating parameters of the hardware allocating machine-learning model **1408** are provided in relation to FIGS. **10A-10B** above.

(203) In addition to FIGS. **1-10B** and FIGS. **12-14**, the corresponding text, and the examples provide a number of different methods, systems, devices, and non-transitory computer-readable media of the intelligent selection and execution platform **102**. In addition to the foregoing, one or more embodiments can also be described in terms of flowcharts comprising acts for accomplishing a particular result, as shown in FIG. **15**. FIG. **15** may be performed with more or fewer acts. Further, the acts may be performed in differing orders. Additionally, the acts described herein may be repeated or performed in parallel with one another or parallel with different instances of the same or similar acts.

(204) As mentioned, FIG. **15** illustrates a flowchart of a series of acts **1500** for selecting a designated hardware environment in accordance with one or more embodiments. While FIG. **15** illustrates acts according to one embodiment, alternative embodiments may omit, add to, reorder, and/or modify any of the acts shown in FIG. **15**. The acts of FIG. **15** can be performed as part of a method. Alternatively, a non-transitory computer-readable medium can comprise instructions that, when executed by one or more processors, cause a computing device to perform the acts of FIG.

15. In some embodiments, a system can perform the acts of FIG. 15.

(205) As shown in FIG. 15, the series of acts **1500** includes an act **1502** of receiving workload data requesting execution of a task using a hardware environment, an act **1504** of extracting workload features defining characteristics of the task, an act **1506** of determining task routing metrics for a plurality of hardware environments hosted in respective network environments, and an act **1108** of selecting a hardware environment for executing the task based on the workload features and the task routing metrics.

(206) In particular, the act **1502** can include receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a hardware environment, the act **1504** can include extracting, from the workload data, workload features defining characteristics of the task, the act **1506** can include determining task routing metrics for a plurality of hardware environments hosted in respective network environments, and the act **1508** can include selecting, from the plurality of hardware environments, a designated hardware environment for executing the task based on the workload features and the task routing metrics.

(207) For example, in one or more embodiments, the series of acts **1500** includes wherein determining task routing metrics for the plurality of hardware environments further comprises determining a hardware state for each hardware environment of the plurality of hardware environments and wherein selecting, from the plurality of hardware environments, the designated hardware environment for executing the task is based in part on the hardware state of the designated hardware environment. In addition, in one or more embodiments, the series of acts **1500** includes wherein determining task routing metrics for the plurality of hardware environments further comprises determining, based on the workload features, a financial cost metric, an execution time metric, an execution cost metric, or a model fit metric for executing the task on each hardware environment of the plurality of hardware environments and wherein selecting, from the plurality of hardware environments, the designated hardware environment for executing the task is based on two or more of the financial cost metric, the execution time metric, the execution cost metric, and the model fit metric.

(208) Further, in one or more embodiments, the series of acts **1500** includes wherein selecting the designated hardware environment further comprises analyzing the workload data and the task routing metrics for the plurality of hardware environments to determine an optimal hardware environment for executing the task from the plurality of hardware environments and selecting the designated hardware environment for executing the task based on determining that the designated hardware environment is the optimal hardware environment from the plurality of hardware environments. Also, in one or more embodiments, the series of acts includes wherein determining that the designated hardware environment is the optimal hardware environment further comprises generating, utilizing a hardware allocating machine-learning model, an optimization metric and determining that the designated hardware environment is the optimal hardware environment based on the optimization metric.

(209) In addition, in one or more embodiments, the series of acts **1500** includes wherein extracting workload features defining characteristics of the task further comprises determining an estimated processing requirement and an estimated storage requirement for executing the task using a hardware environment. Further, in one or more embodiments, the series of acts **1500** includes wherein selecting the designated hardware environment further comprises analyzing the workload data and the task routing metrics for the plurality of hardware environments of the plurality of hardware environments, wherein the plurality of hardware environments comprises one or more hardware environments and one or more third-party hardware environments and selecting the designated hardware environment for executing the task based on analyzing the workload data and the task routing metrics.

(210) Moreover, in one or more embodiments, the series of acts **1500** includes wherein selecting the designated hardware environment further comprises performing probabilistic load balancing of

each hardware environments of the plurality of hardware environments for executing the task and selecting the designated hardware environment based on the probabilistic load balancing. In addition, in one or more embodiments, the series of acts **1500** includes wherein selecting the designated hardware environment further comprises accessing user feedback metrics about executing tasks using one or more hardware environments of the plurality of hardware environments, determining a historical quality metric that indicates how the one or more hardware environments of the plurality of hardware environments will execute the task based on the user feedback metrics, and selecting the designated hardware environment for the task based on the historical quality metric.

(211) Also, in one or more embodiments, the series of acts **1500** includes wherein selecting the designated hardware environment further comprises identifying a capability or specialty associated with one or more hardware environments of the plurality of hardware environments and selecting the designated hardware environment for executing the task based on alignment of the capability or the specialty of the designated hardware environment with the task.

(212) In addition, in one or more embodiments, the series of acts **1500** includes receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a hardware environment, extracting, from the workload data, workload features defining characteristics of the task, determining task routing metrics for a plurality of hardware environments hosted in respective network environments, and selecting, from the plurality of hardware environments and based on the workload features and the task routing metrics, a first designated hardware environment for executing a first portion of the task and a second designated hardware environment for executing a second portion of the task. Further, in one or more embodiments, the series of acts include determining a hardware state for each hardware environment of the plurality of hardware environments, selecting the first designated hardware environment for executing the first portion of the task is based on a first hardware state for the first designated hardware environment and selecting the second designated hardware environment for executing the second portion of the task is based on a second hardware state for the second designated hardware environment.

(213) Further, in one or more embodiments, the series of acts **1500** includes analyzing the workload data and the task routing metrics for the plurality of hardware environments, selecting the first designated hardware environment for executing the first portion of the task based on determining that the first designated hardware environment is a first optimal hardware environment from the plurality of hardware environments, and selecting the second designated hardware environment for executing the second portion of the task based on determining that the second designated hardware environment is a second optimal hardware environment for executing the second portion of the task. Moreover, in one or more embodiments, the series of acts **1500** includes performing probabilistic load balancing of each hardware environment of the plurality of hardware environments for executing the task and selecting the first designated hardware environment for executing the first portion of the task and the second designated hardware environment for executing the second portion of the task based on the probabilistic load balancing.

(214) Also, in one or more embodiments, the series of acts **1500** includes receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a hardware environment, extracting, from the workload data, workload features defining characteristics of the task, determining task routing metrics for a plurality of hardware environments hosted in respective network environments, and selecting, utilizing a hardware allocating machine-learning model, a designated hardware environment from the plurality of hardware environments for executing the task based on the workload features and the task routing metrics.

(215) In addition, in one or more embodiments, the series of acts **1500** includes utilizing the hardware allocating machine-learning model to compare each hardware environment of the

plurality of hardware environments based on the workload features and the task routing metrics and selecting the designated hardware environment from the plurality of hardware environments based on an output of the hardware allocating machine-learning model. Also, in one or more embodiments, the series of acts **1500** includes receiving additional task routing metrics for an additional hardware environment and updating parameters of the hardware allocating machine-learning model based on the additional task routing metrics of the additional hardware environment.

(216) Moreover, in one or more embodiments, the series of acts **1500** includes generating, utilizing the hardware allocating machine-learning model, an optimization metric for each hardware environment of the plurality of hardware environments and selecting the designated hardware environment based on the optimization metric. Further, in one or more embodiments, the series of acts **1500** includes receiving user feedback data indicating a user satisfaction with execution of the task and updating parameters of the hardware allocating machine-learning model based on the user feedback data.

(217) As previously mentioned, the intelligent selection and execution platform **102** can select additional machine-learning models (or hardware environments) as fallbacks for the selected machine-learning models (or hardware environments). In particular, the intelligent selection and execution platform **102** selects a primary machine-learning model (or primary hardware environment) and a fallback machine-learning model (or fallback hardware environment) for executing a task. FIG. **16** illustrates a schematic diagram of an intelligent selection and execution platform selecting a primary machine-learning model (or primary hardware environment) and a fallback machine-learning model (or fallback hardware environment) for executing a task in accordance with one or more embodiments.

(218) As illustrated in FIG. **16**, the intelligent selection and execution platform **102** performs an act **1602** and receives workload data requesting the execution of a task. In particular, the intelligent selection and execution platform **102** receives workload data from a device connected to the content management system **108** that requests the execution of a task. For example, workload data comprises the necessary data and/or information for executing the task, such as a request to execute the task using a machine-learning model. In one or more embodiments, the intelligent selection and execution platform **102** receives workload data through an API layer, as described in relation to FIG. **4** above.

(219) In one or more embodiments, the intelligent selection and execution platform **102** workload data comprises parameters for executing a task. Specifically, the workload data comprises parameters comprising requests or specifications for executing a task. For example, workload data can comprise a parameter that instructs the intelligent selection and execution platform **102** to utilize a certain machine-learning model for executing a task or to execute the task using a certain hardware environment. As another example, workload data can comprise a parameter that instructs the intelligent selection and execution platform **102** to not utilize a certain machine-learning model or hardware environment to execute a task. Additional detail regarding the intelligent selection and execution platform **102** receiving workload data and parameters for executing a task is provided in relation to FIG. **6** above.

(220) As also illustrated in FIG. **16**, the intelligent selection and execution platform **102** performs an act **1604** and extracts, from the workload data, workload features defining characteristics of the task. In particular, the intelligent selection and execution platform **102** extracts workload features by determining an amount of processing and storage needed for executing the task. In some instances, the intelligent selection and execution platform **102** can compute the computational cost for executing a task on a specific machine-learning model. In other instances, the intelligent selection and execution platform **102** determines processing and storage requirements for executing a task using a specific hardware environment. Additional details regarding the intelligent selection and execution platform **102** extracting workload features from workload data are provided in

relation to FIG. 6 above.

(221) As further illustrated in FIG. 16, the intelligent selection and execution platform **102** performs an act **1606** and selects a primary machine-learning model and a fallback machine-learning model. In particular, the intelligent selection and execution platform **102** determines scores based on workload features (e.g., optimization metrics) for machine-learning models and selects a primary machine-learning model and a fallback machine-learning model for executing a task. The intelligent selection and execution platform **102** can then utilize the scores to rank machine-learning models and create a hierarchy (e.g., a daisy chain) of one or more machine-learning models according to the ranking. In some cases, the intelligent selection and execution platform **102** utilizes a model selection machine-learning model to select the primary machine-learning model and fallback machine-learning model. Additional details regarding the intelligent selection and execution platform **102** utilizing a model selection machine-learning model to select a primary machine-learning model and a fallback machine-learning model are provided in relation to FIG. 17 below.

(222) In one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics for a plurality of machine-learning models and utilizes the task routing metrics to select a primary machine-learning model and a fallback machine-learning model. Specifically, the intelligent selection and execution platform **102** determines task routing metrics that indicate financial costs, execution times, execution costs, model fit, model states, model capabilities, or model specialties for executing the task on machine-learning models of a plurality of machine-learning models. Additional details regarding the intelligent selection and execution platform **102** determining task routing metrics for a plurality of machine-learning models are provided in relation to FIG. 7 above. Further, additional details regarding the intelligent selection and execution platform **102** utilizing task routing metrics to select a primary machine-learning model and a fallback machine-learning model are provided in relation to FIG. 17 below.

(223) As also illustrated in FIG. 16, the intelligent selection and execution platform **102** performs an act **1606** and selects a primary hardware environment and a fallback hardware environment. In particular, the intelligent selection and execution platform **102** determines scores based on workload features (e.g., optimization metrics) for hardware environments and selects a primary hardware environment and a fallback hardware environment for executing a task. The intelligent selection and execution platform **102** can then utilize the scores to rank hardware environments and create a hierarchy (e.g., a daisy chain) of two or more hardware environments according to the ranking. In some cases, the intelligent selection and execution platform **102** utilizes a hardware allocating machine-learning model to select the primary hardware environment and the fallback hardware environment. Additional detail regarding the intelligent selection and execution platform **102** utilizing a hardware allocating machine-learning model to select a primary hardware environment and a fallback hardware environment is provided in relation to FIGS. 19A-19B below.

(224) In one or more embodiments, the intelligent selection and execution platform **102** determines task routing metrics for a plurality of hardware environments and utilizes the task routing metrics to select a primary hardware environment and a fallback hardware environment. Specifically, the intelligent selection and execution platform **102** determines task routing metrics that indicate financial costs, execution times, execution costs, user feedback data, hardware environment states, hardware environment capabilities, or hardware environment specialties of hardware environments of a plurality of hardware environments. Additional details regarding the intelligent selection and execution platform **102** determining task routing metrics for a plurality of hardware environments are provided with relation to FIG. 13 above. Further, additional detail regarding the intelligent selection and execution platform **102** utilizing task routing metrics to select a primary hardware environment and a fallback hardware environment is provided with relation to FIGS. 19A-19B below.

(225) As also illustrated in FIG. 16, the intelligent selection and execution platform **102** can

perform an act **1608** and provide data to the fallback machine-learning model to execute a task if the primary machine-learning model is unavailable. In addition, if the fallback machine-learning model is unavailable, the intelligent selection and execution platform **102** continues along the hierarchy and provides data to an additional model to execute the task if the fallback machine-learning model is also unavailable. Additional details regarding the intelligent selection and execution platform **102** providing data to a fallback machine-learning model if a primary machine-learning model is unavailable are provided in relation to FIG. **18** below.

(226) As further illustrated in FIG. **16** as a part of act **1608**, the intelligent selection and execution platform **102** utilizes a fallback hardware system if the primary hardware system is unavailable to execute a task. In addition, if the fallback hardware environment is unavailable, the intelligent selection and execution platform **102** can continue along the hierarchy and utilize an additional hardware environment to execute a task if a fallback hardware environment is unavailable. Additional detail regarding the intelligent selection and execution platform **102** utilizing a fallback hardware environment if a primary hardware environment is unavailable is provided in relation to FIGS. **20A-20B** below.

(227) As previously mentioned, the intelligent selection and execution platform **102** can select a primary machine-learning model and a fallback machine-learning model for executing a task. In particular, the intelligent selection and execution platform **102** utilizes a model selection machine-learning model to select a primary machine-learning model and a fallback machine-learning model based on workload features and task routing metrics. FIG. **17** illustrates the intelligent selection and execution platform **102** utilizing a model selection machine-learning model to select a primary machine-learning model and a fallback machine-learning model in accordance with one or more embodiments.

(228) As illustrated in FIG. **17**, model selection machine-learning model **1716** receives workload features **1702** and task routing metrics **1704**. In particular, the intelligent selection and execution platform **102** extracts workload features **1702** defining characteristics of a task from workload data (e.g., as described in FIG. **6**). Additionally, task routing metrics **1704** comprise task routing metrics for a plurality of machine-learning models that includes a trained machine-learning model **1706** of the intelligent selection and execution platform **102** and third-party machine-learning models in their respective environments (e.g., as described in FIG. **7**). For example, as pictured in FIG. **17**, the plurality of machine-learning models comprises third-party machine-learning model **1710** hosted on third-party server **1708** and third-party machine-learning model **1714** hosted on third-party server **1712** in addition to trained machine-learning model **1706**.

(229) As illustrated, the intelligent selection and execution platform **102** uses model selection machine-learning model **1716** to select a primary machine-learning model **1720**, a fallback machine-learning model **1718**, and an additional machine-learning model **1722**. In particular, model selection machine-learning model **1716** analyzes the workload features **1702** and task routing metrics **1704** to select primary machine-learning model **1720**, fallback machine-learning model **1718**, and additional machine-learning model **1722**. For example, the model selection machine-learning model **1716** compares estimated processing requirements and estimated storage requirements for executing the task to the task routing metrics **1704** for the plurality of machine-learning models. For example, the model selection machine-learning model **1716** determines that the execution cost metric or the model state aligns with the estimated processing requirement or the estimated storage requirement for executing the task. As another example, if a model capability or model specialty of a machine-learning model aligns with a task, the model selection machine-learning model **1716** can select that machine-learning model for executing the task. Moreover, as another example, the model selection machine-learning model **1716** can select primary machine-learning model **1720**, fallback machine-learning model **1718**, and additional machine-learning model **1722** based on a financial cost metric, such as by ranking machine-learning models of the plurality of machine-learning models based on a financial cost metric.

(230) In one or more embodiments, the model selection machine-learning model **1716** generates an optimization metric based on workload features **1702** and task routing metrics **1704** and uses the optimization metric to select primary machine-learning model **1720**, fallback machine-learning model **1718**, and additional machine-learning model **1722**. Specifically, the model selection machine-learning model **1716** generates an optimization metric for each machine-learning model of the plurality of machine-learning models and ranks the machine-learning models based on the optimization metric to generate a hierarchy of machine-learning models for executing a task. In some cases, the model selection machine-learning model **1716** selects primary machine-learning model **1720**, fallback machine-learning model **1718**, and additional machine-learning model **1722** based on the hierarchy (e.g., primary machine-learning model **1720** has the highest optimization metric, fallback machine-learning model **1718** has the second highest optimization metric, and additional machine-learning model **1722** has a third highest optimization metric).

(231) As previously mentioned, the intelligent selection and execution platform **102** provides data to a fallback machine-learning model if a primary machine-learning model is unavailable. In particular, the intelligent selection and execution platform **102** provides data to a fallback machine-learning model to execute a task if a primary machine-learning model is unavailable. FIG. **18** illustrates an intelligent selection and execution platform utilizing a fallback machine-learning model to execute a task when a primary machine-learning model is unavailable in accordance with one or more embodiments.

(232) As illustrated in FIG. **18**, model selection machine-learning model **1802** provides the output of designated machine-learning models **1804** comprising a primary machine-learning model, a fallback machine-learning model, and an additional model. In particular, as described in FIG. **17** above, model selection machine-learning model **1802** generates a hierarchy of machine-learning models based on workload features defining characteristics of a task and task routing metrics of a plurality of machine-learning models.

(233) As also illustrated in FIG. **18**, the intelligent selection and execution platform **102** provides data to primary machine-learning model **1806** to execute a task and determines that the primary hardware environment is unavailable. In some cases, the intelligent selection and execution platform **102** attempts to execute the task using primary machine-learning model **1806** and determines that it is unavailable. In other cases, the intelligent selection and execution platform **102** identifies that primary machine-learning model **1806** is unavailable prior to attempting to execute the task.

(234) Primary machine-learning model **1806**, or any of the designated machine-learning models **1804**, may be unavailable for a number of reasons. For example, primary machine-learning model **1806** may be unavailable due to an outage associated with the primary machine-learning model. As another example, hardware, systems, or network environments associated with primary machine-learning model **1806** may experience issues, such as power outages or other technical issues, rendering primary machine-learning model **1806** unavailable. In addition, as another example, primary machine-learning model **1806** is unavailable due to high demand, such as a high volume of requests to execute tasks. Further, as another example, primary machine-learning model **1806** may be unavailable due to software changes or updates. Moreover, as another example, primary machine-learning model **1806** may be unavailable due to natural disasters that damage hardware supporting primary machine-learning model **1806** or restrict communication with primary machine-learning model **1806**.

(235) As illustrated in FIG. **18**, upon detecting that primary machine-learning model **1806** is unavailable, the intelligent selection and execution platform **102** provides data to fallback machine-learning model **1808** to execute task **1810**. Moreover, as illustrated in FIG. **16**, if the intelligent selection and execution platform **102** detects that fallback machine-learning model **1808** is unavailable, the intelligent selection and execution platform **102** can provide data to additional machine-learning model **1812** to execute task **1810**. Indeed, the intelligent selection and execution

platform **102** can continue to move along a hierarchy of machine-learning models until a machine-learning model is available to execute task **1810**.

(236) As also illustrated in FIG. **18**, the intelligent selection and execution platform **102** receives user feedback data **1814**. In particular, as previously mentioned, the intelligent selection and execution platform **102** receives user feedback data **1814** that indicates a user satisfaction with a task and utilizes the user feedback data **1814** to update parameters of model selection machine-learning model **1802**. In some cases, the intelligent selection and execution platform **102** accounts for model unavailability when utilizing user feedback data **1814** to update parameters of model selection machine-learning model **1802**. For example, if a client device indicated dissatisfaction with a task, the intelligent selection and execution platform **102** might indicate that a preferred machine-learning model was unavailable to execute the task. The intelligent selection and execution platform **102** may also account for model availability when updating parameters of model selection machine-learning model **1802** with user feedback data **1814**.

(237) As previously mentioned, intelligent selection and execution platform **102** selects a primary hardware environment and a fallback hardware environment for executing a task. In particular, the intelligent selection and execution platform **102** utilizes a hardware allocating machine-learning model to select a primary hardware environment and a fallback hardware environment for executing a task. FIGS. **19A-19B** illustrate an intelligent selection and execution platform utilizing a hardware allocating machine-learning model to select a primary hardware environment and a fallback hardware environment in accordance with one or more embodiments. FIG. **19A** illustrates a hardware allocating machine-learning model selecting a primary hardware environment and a fallback hardware environment for executing a task. FIG. **19B** illustrates the hardware allocating machine-learning model allocating tasks between multiple primary hardware environments and selects a fallback hardware environment if one of the primary hardware environments is unavailable.

(238) As illustrated in FIG. **19A**, hardware allocating machine-learning model **1916** receives workload features **1902** and task routing metrics **1904**. In particular, the intelligent selection and execution platform **102** extracts workload features **1902** defining characteristic of a task from workload data (e.g., as described in FIG. **6**). Additionally, task routing metrics **1904** comprise task routing metrics for a plurality of hardware environments that includes a hardware environment **1906** of intelligent selection and execution platform **102** and third-party hardware environments in their respective network environments (e.g., as described in FIG. **13**). For example, as illustrated in FIG. **19A**, the plurality of hardware environment comprises third-party hardware environment **1910** hosted on third-party server **1908** and third-party hardware environment **1914** hosted on third-party server **1912** in addition to hardware environment **1906**.

(239) As also illustrated in FIG. **19A**, the intelligent selection and execution platform **102** uses hardware allocating machine-learning model **1916** to select a primary hardware environment **1920**, a fallback hardware environment **1918**, and an additional hardware environment **1922**. In particular, hardware allocating machine-learning model **1916** analyzes workload features **1902** and task routing metrics **1904** to select primary hardware environment **1920**, fallback hardware environment **1918**, and additional hardware environment **1922**. For example, hardware allocating machine-learning model **1916** compares estimated processing requirements and storage requirements for executing the task to the task routing metrics for each hardware environment of the plurality of hardware environments. For example, the hardware allocating machine-learning model **1916** determines that execution cost metrics or model state metrics align with the estimated processing requirement or the estimated storage requirement for executing the task. As another example, if a hardware environment capability or hardware environment specialty of a hardware environment aligns with a task, hardware allocating machine-learning model **1916** can select that hardware environment for executing the task. Moreover, as another example, the hardware allocating machine-learning model **1916** can select primary hardware environment **1920**, fallback

hardware environment **1918**, and additional hardware environment **1922** based on a financial cost metric, such as by ranking machine-learning models of the plurality of machine-learning models based on a financial cost metric.

(240) In one or more embodiments, hardware allocating machine-learning model **1916** generates an optimization metric based on workload features **1902** and task routing metrics **1904** and utilizes the optimization metric to select primary hardware environment **1920**, fallback hardware environment **1918**, and additional hardware environment **1922**. Specifically, hardware allocating machine-learning model **1916** generates an optimization metric for each hardware environment of the plurality of hardware environments and ranks the hardware environments based on the optimization metric to generate a hierarchy of hardware environments for executing a task. In some cases, the hardware allocating machine-learning model **1916** selects select primary hardware environment **1920**, fallback hardware environment **1918**, and additional hardware environment **1922** based on the hierarchy (e.g., primary hardware environment **1920** has the highest optimization metric, fallback hardware environment **1918** has the second highest optimization metric, and additional hardware environment **1922** has a third highest optimization metric).

(241) In one or more embodiments, the intelligent selection and execution platform **102** selects a primary hardware environment by allocating resources from multiple hardware environments to execute a task. As shown in FIG. **19B**, hardware allocating machine-learning model **1916** can select hardware environment **1906** as a first primary hardware environment **1924** and third-party hardware environment **1910** as a second primary hardware environment **1926** and allocate from both the first primary hardware environment **1924** and the second primary hardware environment **1926** to execute the task. For example, as illustrated, the hardware allocating machine-learning model allocates a first portion of a task (e.g., 50%) to the first primary hardware environment **1924** and a second portion of a task (e.g., 50%) to the second primary hardware environment **1926**.

(242) As also illustrated in FIG. **19B**, in cases where the intelligent selection and execution platform **102** allocates resources to multiple hardware environments, the intelligent selection and execution platform **102** can also select an additional hardware environment as a fallback hardware environment to execute a task if the first primary hardware environment or the second primary hardware environment, or both, are unavailable. For example, as illustrated, the hardware allocating machine-learning model **1916** can select third-party hardware environment **1914** as fallback hardware environment **1928**.

(243) As previously mentioned, the intelligent selection and execution platform **102** can allocate resources from a fallback hardware environment if a primary hardware environment is unavailable. In particular, the intelligent selection and execution platform **102** allocates resources to a fallback hardware environment to execute a task if a primary hardware environment is unavailable. FIGS. **20A-20B** illustrate an intelligent selection and execution platform detecting a primary hardware environment is unavailable and utilizing a fallback hardware environment to execute a task in accordance with one or more embodiments. FIG. **20A** illustrates the intelligent selection and execution platform **102** allocating resources if a primary hardware environment is unavailable. FIG. **20B** illustrates the intelligent selection and execution platform **102** adjusting the distribution of resources if a hardware environment is unavailable.

(244) As illustrated in FIG. **20A**, hardware allocating machine-learning model **2002** provides an output of designated hardware environments **2004** comprising a primary hardware environment, a fallback hardware environment, and additional hardware environments. In particular, as described in FIG. **19A** above, hardware allocating machine-learning model **2002** generates a hierarchy of hardware environments based on workload features defining characteristics of a task and task routing metrics of a plurality of hardware environments. In some instances, as described in FIG. **19B** above, the hardware allocating machine-learning model can allocate execution of tasks between multiple hardware environments.

(245) As also illustrated in FIG. **20A**, the intelligent selection and execution platform **102** utilizes

primary hardware environment **2006** to execute a task and determines that primary hardware environment **2006** is unavailable. In some cases, the intelligent selection and execution platform **102** attempts to utilize primary hardware environment **2006** and determines that it is unavailable. In other cases, the intelligent selection and execution platform **102** identifies that primary hardware environment **2006** is unavailable prior to attempting to execute the task.

(246) Primary hardware environment **2006**, or another designated hardware environment of designated hardware environments **2004**, may be unavailable for a number of reasons. For example, primary hardware environment **2006** may be unavailable due to hardware failure, such as when a GPU, CPU, memory, or storage drives fail. As another example, primary hardware environment **2006** may be unavailable due to resource constraints, such as when computational resources are overcommitted to other tasks. As an additional example, the primary hardware environment **2006** may be unavailable due to scheduled maintenance or upgrades that temporarily make the hardware environment unavailable. Moreover, for example, primary hardware environment **2006** may be unavailable due to power outages, power disruptions or other natural factors. In addition, for example, primary hardware environment **2006** may be unavailable due to cooling issues, such as when there are issues with cooling systems, and primary hardware environment **2006** may be shut down to prevent overheating. As another example, primary hardware environment **2006** may be unavailable due to network connectivity problems that prevent primary hardware environment **2006** from communicating with machine-learning models or other systems. Further, for example, primary hardware environment **2006** may be unavailable due to unexpected demand spikes, such as a sudden increase in demand for computational resources due to unexpected user activity or other factors. Moreover, primary hardware environment **2006** may be unavailable due to errors, such as software or configuration errors that prevent communication or misallocate resources, leading to misallocation of resources.

(247) As illustrated in FIG. **20A**, upon detecting that primary hardware environment **2006** is unavailable, the intelligent selection and execution platform **102** utilizes fallback hardware environment **2008** to execute task **2012**. Moreover, as illustrated in FIG. **20A**, if the intelligent selection and execution platform **102** detects that fallback hardware environment **2008** is unavailable, the intelligent selection and execution platform **102** can utilize additional hardware environment **2010** to execute task **2012**. Indeed, the intelligent selection and execution platform **102** can continue to move along a hierarchy of hardware environments until a hardware environment is available to provide computational resources for executing task **2012**.

(248) As also illustrated in FIG. **20A**, the intelligent selection and execution platform **102** receives user feedback data **2014**. In particular, as previously mentioned, the intelligent selection and execution platform **102** receives user feedback data **2014** that indicates a user satisfaction with a task and utilizes user feedback data **2014** to update parameters of hardware allocating machine-learning model **2002**. In some cases, the intelligent selection and execution platform **102** accounts for hardware environment availability when utilizing user feedback data **2014** to update parameters of hardware allocating machine-learning model **2002**. For example, if a client device indicated dissatisfaction with a task, the intelligent selection and execution platform **102** might indicate that a preferred machine-learning model was unavailable to execute the task. The intelligent selection and execution platform **102** may also account for model availability when updating parameters of hardware allocating machine-learning model **2002**. Additional detail regarding the intelligent selection and execution platform **102** utilizing user feedback data to update parameters of a hardware allocating machine-learning model is provided in relation to FIGS. **10A-10B** above.

(249) As previously mentioned, the intelligent selection and execution platform **102** can allocate computational resources from multiple hardware environments for executing a task. As illustrated in FIG. **20B**, the intelligent selection and execution platform **102** may redistribute or reallocate computational resources if a primary hardware environment is unavailable. As shown, hardware allocating machine-learning model **2002** can select hardware environment **2016** of the intelligent

selection and execution platform **102** as a first primary hardware environment, third-party hardware environment **2020** hosted on third-party server **2018** as a second primary hardware environment, and third-party hardware environment **2024** hosted on third-party server **2022** as a fallback hardware environment. If the intelligent selection and execution platform **102** determines that hardware environment **2016** or third-party hardware environment **2020** are unavailable, the intelligent selection and execution platform **102** can allocate resources from third-party hardware environment **2024** (e.g., the fallback hardware environment). As shown, the intelligent selection and execution platform **102** allocates resources from hardware environment **2016** to execute a first portion of a task (50%) and allocates resources from third-party hardware environment **2020** to execute a second portion of a task (50%). Upon detecting that third-party hardware environment **2020** is unavailable, the intelligent selection and execution platform **102** allocates resources to the third-party hardware environment **2024** to execute the second portion of the task.

(250) In addition to FIGS. **1-10B**, FIGS. **12-14**, and FIGS. **16-20B**, the corresponding text, and the examples provide a number of different methods, systems, devices, and non-transitory computer-readable media of the intelligent selection and execution platform **102**. In addition to the foregoing, one or more embodiments can also be described in terms of flowcharts comprising acts for accomplishing a particular result, as shown in FIG. **21**. FIG. **21** may be performed with more or fewer acts. Further, the acts may be performed in differing orders. Additionally, the acts described herein may be repeated or performed in parallel with one another or parallel with different instances of the same or similar acts.

(251) As mentioned, FIG. **21** illustrates a flowchart of a series of acts **2100** for selecting a machine-learning model and a fallback machine-learning model in accordance with one or more embodiments. While FIG. **21** illustrates acts according to one embodiment, alternative embodiments may omit, add to, reorder, and/or modify any of the acts shown in FIG. **21**. The acts of FIG. **21** can be performed as part of a method. Alternatively, a non-transitory computer-readable medium can comprise instructions that, when executed by one or more processors, cause a computing device to perform the acts of FIG. **21**. In some embodiments, a system can perform the acts of FIG. **21**.

(252) As shown in FIG. **21**, the series of acts **2100** includes an act **2102** of receiving workload data requesting execution of a task using a machine-learning model, an act **2104** of extracting workload features defining characteristics of the task, an act **2106** of selecting a primary machine-learning model and a fallback machine-learning model for executing the task based on the workload features, and an act **2108** of based on detecting that the primary machine-learning model is unavailable, providing the workload data to a computing environment of the fallback machine-learning model for executing the task.

(253) In particular, the act **2102** can include receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model, the act **2104** can include extracting, from the workload data, workload features defining characteristics of the task, the act **2106** can include selecting a primary machine-learning model and a fallback machine-learning model for executing the task based on the workload features and the act **2108** can include based on detecting that the primary machine-learning model is unavailable, providing the workload data to a computing environment of the fallback machine-learning model for executing the task.

(254) For example, in one or more embodiments, the series of acts **2100** includes wherein extracting workload features defining characteristics of the task further comprises determining an estimated processing requirement and an estimated storage requirement for executing the task. In addition, in one or more embodiments, the series of acts **2100** includes wherein selecting the primary machine-learning model and the fallback machine-learning model further comprises generating optimization metrics for each machine-learning model of a plurality of machine-learning model and selecting the primary machine-learning model and the fallback machine-

learning model based on the optimization metrics.

(255) Further, in one or more embodiments, the series of acts **2100** includes wherein selecting the primary machine-learning model and the fallback machine-learning model further comprises determining a model state for a plurality of machine-learning models, selecting the primary machine-learning model based in part on a first model state of the primary machine-learning model, and selecting the fallback machine-learning model based in part on a second model state of the fallback machine-learning model. Also, in one or more embodiments, the series of acts **2100** includes wherein selecting the primary machine-learning model and the fallback machine-learning model further comprises determining, based on the workload features, a financial cost metric, an execution time metric, an execution cost metric, or a model fit metric for executing the task on each machine-learning models in a plurality of machine-learning models and selecting the primary machine-learning model and the fallback machine-learning model is based on two or more of the financial cost metric, the execution time metric, the execution cost metric, or the model fit metric

(256) In addition, in one or more embodiments, the series of acts **2100** includes selecting an additional machine-learning model for executing the task based on the workload feature and, based on detecting that the fallback machine-learning model is unavailable, providing the workload data to a computing environment of the additional machine-learning model for executing the task. Moreover, the series of acts **2100** includes selecting a trained machine-learning model as the primary machine-learning model and a third-party trained machine-learning model as the fallback machine-learning model.

(257) Also, in one or more embodiments, the series of acts **2100** includes detecting that the primary machine-learning model is unavailable based on identifying that a first hardware environment associated with the primary machine-learning model is unavailable to execute the task and providing the workload data to a second hardware environment associated with the fallback machine-learning model based on identifying that the second hardware environment is available to execute the task. In addition, in one or more embodiments, the series of acts **2100** includes wherein selecting the primary machine-learning model and the fallback machine-learning model further comprises performing a software domain analysis of each machine-learning model of a plurality of machine-learning models for executing the task and selecting the primary machine-learning model and the fallback machine-learning model based on the software domain analysis.

(258) Moreover, in one or more embodiments, the series of acts **2100** includes receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model, extracting, from the workload data, workload features defining characteristics of the task, determining task routing metrics for a plurality of machine-learning models hosted in respective network environments, selecting a primary machine-learning model and a fallback machine-learning model from the plurality of machine-learning models for executing the task based on the workload features and the task routing metrics, and, based on detecting that the primary machine-learning model is unavailable, provide the workload data to a computing environment of the fallback machine-learning model for executing the task.

(259) In addition, in one or more embodiments, the series of acts **2100** includes selecting the primary machine-learning model and the fallback machine-learning model by determining task routing metrics for the plurality of machine-learning models by determining a financial cost metric, an execution time metric, an execution cost metric, or a model fit metric for executing the task on each machine-learning models of the plurality of machine-learning models and selecting, from the plurality of machine-learning models, the primary machine-learning model and the fallback machine-learning model based on two or more of the financial cost metric, the execution time metric, the execution cost metric, or the model fit metric for executing the task. Also, in one or more embodiments, the series of acts **2100** includes selecting the primary machine-learning model and the fallback machine-learning model by generating, based on the workload data and the task routing metrics, an optimization metric for each machine-learning model of the plurality of

machine-learning models and selecting the primary machine-learning model and the fallback machine-learning model from the plurality of machine-learning models based on a first optimization metric for the primary machine-learning model and a second optimization metric for the fallback machine-learning model.

(260) Further, in one or more embodiments, the series of acts **2100** includes selecting the primary machine-learning model and the fallback machine-learning model by accessing historical user feedback data about executing tasks using one or more machine-learning models of the plurality of machine-learning models, determining a historical quality metric for each machine-learning model of the plurality of machine-learning models based on the historical user feedback data, and selecting the primary machine-learning model and the fallback machine-learning model from the plurality of machine-learning models based in part on a first historical quality metric for the primary machine-learning model and a second historical quality metric for the fallback machine-learning model. In addition, in one or more embodiments, the series of acts **2100** includes select the primary machine-learning model and the fallback machine-learning model by identifying a capability or a specialty associated with one or more machine-learning models of the plurality of machine-learning models and selecting the primary machine-learning model for executing the task based on alignment of the capability or the specialty of the primary machine-learning model with the task. Moreover, in one or more embodiments, the series of acts **2100** includes selecting a trained machine-learning model as the primary machine-learning model and a third-party trained machine-learning model as the fallback machine-learning model.

(261) Moreover, in one or more embodiments, the series of acts **2100** includes receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model, extracting, from the workload data, workload features defining characteristics of the task, selecting, utilizing a model selection machine-learning model, a primary machine-learning model and a fallback machine-learning model for executing the task based on the workload features, and, based on detecting that the primary machine-learning model is unavailable, provide the workload data to a computing environment of the fallback machine-learning model for executing the task.

(262) Also, in one or more embodiments, the series of acts **2100** includes utilizing the model selection machine-learning model to compare each machine-learning model in a plurality of machine-learning models based on the workload features, where the plurality of machine-learning models comprises one or more trained models and one or more third-party trained models and selecting the primary machine-learning model and the fallback machine-learning model from the plurality of machine-learning models based on an output of the model selection machine-learning model. In addition, in one or more embodiments, the series of acts **2100** acts includes training the model selection machine-learning model based on task routing metrics of a plurality of machine-learning models, where the plurality of machine-learning models comprises the primary machine-learning model and the fallback machine-learning model, receiving additional task routing metrics for an additional machine-learning model, and updating parameters of the model selection machine-learning model based on the additional task routing metrics of the additional machine-learning model.

(263) Further, in one or more embodiments, the series of acts **2100** includes receiving user feedback data indicating a user satisfaction with performance of the primary machine-learning model and updating parameters of the model selection machine-learning model based on the user feedback data. Moreover, in one or more embodiments, the series of acts **2100** includes generating, utilizing the model selection machine-learning model, an optimization metric for each machine-learning model of a plurality of machine-learning models and selecting the primary machine-learning model and the fallback machine-learning model from the plurality of machine-learning models based on a first optimization metric for the primary machine-learning model and a second optimization metric for the fallback machine-learning model.

(264) As previously mentioned, the intelligent selection and execution platform **102** can execute training tasks. In particular, the intelligent selection and execution platform **102** trains machine-learning models for their respective tasks using particular hardware environments and based on bandwidth availability. FIG. 22 illustrates a schematic diagram of an intelligent selection and execution platform **102** intelligently initiating and pausing training tasks based in accordance with one or more embodiments.

(265) As illustrated, the intelligent selection and execution platform **102** performs an act **2202** and monitors hardware usage metrics for a hardware environment. For example, the intelligent selection and execution platform **102** can monitor hardware usage metrics by measuring, determining, or receiving metrics indicating the usage of GPU, CPU, memory, or storage components of the hardware environment. In addition, the intelligent selection and execution platform **102** measures, determines or receives bandwidth metrics of the hardware environment, such as throughput, latency, jitter, or packet loss. Moreover, in addition to measuring hardware usage metrics of the hardware environment, the intelligent selection and execution platform **102** can utilize the task routing metrics as described above in relation to FIG. 6. For instance, the intelligent selection and execution platform **102** can utilize a hardware environment state that indicates bandwidth availability for the hardware environment.

(266) As further illustrated in FIG. 22, the intelligent selection and execution platform **102** performs an act **2204** and receives workload data requesting the execution of a training task. Specifically, the workload data comprises information necessary to execute the training task. For example, the workload data indicates whether the training task is a hot path task that requires continuous (unbroken) execution by a model/hardware (e.g., pausing or delaying the task could result in performance degradation or service disruption. As another example, the intelligent selection and execution platform **102** receives workload data indicating an anticipated amount of time and processing capability required for the training task. Additional details regarding the intelligent selection and execution platform **102** receiving workload data are provided in relation to FIG. 6 above.

(267) In one or more embodiments, the intelligent selection and execution platform **102** can also receive workload data requesting execution of a batch task. In particular, the intelligent selection and execution platform **102** receives workload data requesting the execution of a batch task comprising a set or grouping of tasks that are queued for execution together. For example, the intelligent selection and execution platform **102** can execute portions of the batch task a

(268) As further illustrated in FIG. 22, the intelligent selection and execution platform **102** performs an act **2206** and initiates the training task based on bandwidth availability. In particular, the intelligent selection and execution platform **102** determines points when the hardware usage metrics indicate a minimum-use time period (e.g., a “trough”) where the hardware environment has computing resources for a training task (or batch task) and initiates the training task or batch task. For example, the intelligent selection and execution platform **102** can determine, based on the hardware usage metrics, that there is availability (e.g., idle GPUs or CPUs) in a hardware environment within the intelligent selection and execution platform **102** or in a third-party hardware environment and initiate the training task (or batch task).

(269) Moreover, in addition to initiating a training task, the intelligent selection and execution platform **102** may also schedule training tasks based on bandwidth availability. In particular, the intelligent selection and execution platform **102** can monitor bandwidth availability throughout a use time period to identify minimum-use time periods and high-use time periods and schedule training tasks (or batch tasks) during minimum-use time periods. For example, the intelligent selection and execution platform **102** may forecast a minimum-use time period, where an average bandwidth availability satisfies a training task bandwidth usage threshold and schedule the training task (or batch task) during the minimum-use time period. As another example, the intelligent selection and execution platform **102** can forecast a high-use time period, where an average

bandwidth availability satisfies a training task bandwidth usage threshold (or a batch task bandwidth usage threshold) and determines not to schedule a training task (or batch task) during the high-use time period. Additional detail regarding the intelligent selection and execution platform **102** initiating or scheduling training (or batch) tasks based on bandwidth availability is provided in relation to FIG. **23** below.

(270) As illustrated in FIG. **22**, the intelligent selection and execution platform **102** can perform an act **2208** and pause the training task based on detecting a change in bandwidth availability. In particular, the intelligent selection and execution platform **102** can pause the training task (or batch task) based on detecting that bandwidth availability decreased (e.g., other tasks are now using previously free bandwidth). For example, the intelligent selection and execution platform **102** can determine to pause a training task (or batch task) based on a training task bandwidth usage threshold (or batch task bandwidth usage threshold). In some instances, the intelligent selection and execution platform **102** will pause the training task based on detecting that bandwidth no longer satisfies a training task bandwidth usage threshold (or batch task bandwidth usage threshold). In addition, the intelligent selection and execution platform **102** can pause a training task (or batch task) based on detecting a hot path task that needs execution (e.g., hot path tasks take priority). Additional details regarding the intelligent selection and execution platform **102** pausing training tasks are provided in relation to FIG. **23** below.

(271) As mentioned, the intelligent selection and execution platform **102** initiates and pauses training tasks. In particular, the intelligent selection and execution platform **102** initiates training tasks based on bandwidth availability and pauses training tasks based on bandwidth availability. FIG. **23** illustrates an example graph depicting bandwidth availability based on usage throughout time in accordance with one or more embodiments.

(272) As illustrated in FIG. **23**, the intelligent selection and execution platform **102** can determine bandwidth availability **2302** over a use time period **2304**. In particular, the intelligent selection and execution platform **102** determines bandwidth availability **2302** during a use time period **2304** to determine minimum-use time periods and high-use time periods. For example, as illustrated, the intelligent selection and execution platform **102** can identify a minimum-use time period **2310** or a minimum-use time period **2314** where bandwidth availability **2302** indicates a lower bandwidth usage and a high-use time period **2312** where bandwidth availability **2302** indicates a higher bandwidth usage.

(273) In one or more embodiments, the intelligent selection and execution platform **102** utilizes a threshold **2306** to identify a minimum-use time period **2310** or a high-use time period **2312**. In particular, when bandwidth availability **2302** satisfies threshold **2306**, the intelligent selection and execution platform **102** may determine there is a minimum-use time period **2310** or a minimum-use time period **2314**. Similarly, when bandwidth availability **2302** no longer satisfies threshold **2306**, the intelligent selection and execution platform **102** may determine there is a high-use time period.

(274) As previously mentioned, the intelligent selection and execution platform **102** initiates and pauses training tasks based on bandwidth availability. In particular, the intelligent selection and execution platform **102** initiates and pauses tasks based on threshold **2306** for the respective task. For example, the intelligent selection and execution platform **102** can initiate a training task based on bandwidth availability **2302**, satisfying a training task bandwidth usage threshold, and pause the training task based on bandwidth availability no longer satisfying the training task bandwidth usage threshold. As another example, the intelligent selection and execution platform **102** can initiate a batch task based on detecting that bandwidth availability **2302** satisfies a batch task bandwidth usage threshold and pausing the batch task based on detecting that bandwidth availability **2302** no longer satisfies the batch task bandwidth usage threshold.

(275) In one or more embodiments, threshold **2306** varies based on the type of task. In particular, the intelligent selection and execution platform **102** utilizes thresholds of different values for different types of tasks. For example, the intelligent selection and execution platform **102** may have

different values for a training task bandwidth usage threshold and a batch task bandwidth usage threshold. As another example, the intelligent selection and execution platform **102** can identify that a training task requires continuous execution or that the estimated training task length is longer than an average minimum-use time period and adjust threshold **2306** (e.g., so it will not pause as easily based on bandwidth availability).

(276) As previously mentioned, the intelligent selection and execution platform **102** can also schedule training tasks and batch tasks. In particular, the intelligent selection and execution platform **102** can schedule a training task for an average minimum-use time period **2310** where the average bandwidth availability **2302** satisfies threshold **2306**. For example, the intelligent selection and execution platform **102** can schedule a training task for minimum-use time period **2310** where bandwidth availability **2302** generally satisfies a training task bandwidth usage threshold. As another example, the intelligent selection and execution platform **102** can schedule a batch task for a minimum-use time period **2314**, where bandwidth availability **2302** generally satisfies a batch task bandwidth usage threshold.

(277) In one or more embodiments, the intelligent selection and execution platform **102** can pause scheduled tasks based on detecting a change in bandwidth availability **2302**. Specifically, the intelligent selection and execution platform **102** pauses scheduled tasks based on determining that bandwidth availability **2302** no longer satisfies threshold **2306**. For example, the intelligent selection and execution platform **102** can pause a scheduled training task based on detecting that bandwidth availability **2302** no longer satisfies a training task bandwidth usage threshold. As another example, the intelligent selection and execution platform **102** can pause a scheduled batch task based on detecting that bandwidth availability **2302** no longer satisfies a batch job bandwidth usage threshold. Moreover, if the intelligent selection and execution platform **102** detects an additional change in bandwidth availability, the intelligent selection and execution platform **102** can initiate the task once bandwidth availability **2302** satisfies threshold **2306** again.

(278) In addition to FIGS. **1-10B**, FIGS. **12-14**, FIGS. **16-20B**, and FIGS. **22-23**, the corresponding text, and the examples provide a number of different methods, systems, devices, and non-transitory computer-readable media of the intelligent selection and execution platform **102**. In addition to the foregoing, one or more embodiments can also be described in terms of flowcharts comprising acts for accomplishing a particular result, as shown in FIG. **24**. FIG. **24** may be performed with more or fewer acts. Further, the acts may be performed in differing orders. Additionally, the acts described herein may be repeated or performed in parallel with one another or parallel with different instances of the same or similar acts.

(279) As mentioned, FIG. **24** illustrates a flowchart of a series of acts **2400** for initiating and pausing training tasks in accordance with one or more embodiments. While FIG. **24** illustrates acts according to one embodiment, alternative embodiments may omit, add to, reorder, and/or modify any of the acts shown in FIG. **24**. The acts of FIG. **24** can be performed as part of a method. Alternatively, a non-transitory computer-readable medium can comprise instructions that, when executed by one or more processors, cause a computing device to perform the acts of FIG. **24**. In some embodiments, a system can perform the acts of FIG. **24**.

(280) As shown in FIG. **24**, the series of acts **2400** includes an act **2402** of monitoring hardware usage metrics for a hardware environment hosting a machine-learning model, an act **2404** of extracting workload data requesting execution of a training task using the machine-learning model, an act **2406** of initiating the training task at the hardware environment based on detecting a bandwidth availability indicated by the hardware usage metrics, and an act **2408** of pausing the training task based on detecting a change in the bandwidth availability indicated by the hardware usage metrics.

(281) In particular, the act **2402** can include monitoring hardware usage metrics for a hardware environment hosting a machine-learning model, the act **2404** can include extracting, at a content management system, workload data requesting execution of a training task using the machine-

learning model, the act **2406** can include based on the workload data, initiating the training task at the hardware environment based on detecting a bandwidth availability indicated by the hardware usage metrics, and the act **2408** can include pausing the training task based on detecting a change in the bandwidth availability indicated by the hardware usage metrics.

(282) For example, in one or more embodiments, the series of acts **2400** includes wherein initiating the training task at the hardware environment further comprises identifying that the hardware usage metrics indicate that the bandwidth availability satisfies a training task bandwidth usage threshold and initiating the training task based on identifying that the bandwidth availability satisfies the training task bandwidth usage threshold. In addition, in one or more embodiments, the series of acts **2400** includes forecasting a minimum-use time period where an average bandwidth availability satisfies a training task bandwidth usage threshold and scheduling the training task for the minimum-use time period.

(283) Also, in one or more embodiments, the series of acts **2400** includes initiating the training task at the hardware environment during the minimum-use time period based on detecting that an actual bandwidth usage at the minimum-use time period satisfies a training task bandwidth usage threshold. Further, in one or more embodiments, the series of acts **2400** includes wherein initiating the training task at the hardware environment further comprises monitoring hardware usage metrics by monitoring a user count metric, a memory usage metric, a central processing unit metric, or a graphics processing unit metric and initiating the training task at the hardware environment based on two or more of the user count metric, the memory usage metric, the central processing unit metric, and the graphics processing unit metric.

(284) In addition, in one or more embodiments, the series of acts **2400** includes wherein pausing the training task further comprises monitoring the hardware usage metrics for the hardware environment throughout a use time period to identify a high-use time period and pausing the training task based on forecasting the change in the bandwidth availability at the high-use time period. Also, in one or more embodiments, the series of acts **2400** includes wherein extracting workload data requesting execution of the training task further comprises determining an estimated processing requirement and an estimated storage requirement for executing the training task and determining an estimated execution length for the training task.

(285) Further, in one or more embodiments, the series of acts **2400** includes receiving the workload data by receiving an indication that the training task does not require continuous execution and initiating the training task based on the indication that the training task does not require continuous executions. In addition, in one or more embodiments, the series of acts **2400** includes initiating the training task based on identifying that the hardware usage metrics indicate that the bandwidth availability satisfies a training task bandwidth usage threshold and pausing the training task based on identifying that the hardware usage metrics indicate that the bandwidth availability no longer satisfies the training task bandwidth usage threshold.

(286) Moreover, in one or more embodiments, the series of acts **2400** includes monitoring hardware usage metrics for a hardware environment hosting a machine-learning model, extracting, at a content management system, workload data requesting execution of a batch task using the machine-learning model, based on the workload data, initiating the batch task at the hardware environment based on detecting a bandwidth availability indicated by the hardware usage metrics, and pausing the batch task based on detecting a change in the bandwidth availability indicated by the hardware usage metrics.

(287) In addition, in one or more embodiments, the series of acts **2400** includes forecasting a minimum-use time period where an average bandwidth availability satisfies a batch task bandwidth usage threshold and scheduling the batch task for the minimum-use time period. Also, in one or more embodiments, the series of acts **2400** includes resuming the batch task based on detecting an additional change in the bandwidth availability. Further, in one or more embodiments, the series of acts **2400** includes extracting workload data by determining an estimated processing requirement

and an estimated storage requirement for executing the batch task and determining an estimated execution length for the batch task.

(288) Also, in one or more embodiments, the series of acts **2400** includes initiating the batch task at the hardware environment by identifying that the hardware usage metrics indicate that the bandwidth availability satisfies a batch task bandwidth usage threshold and initiating the batch task based on identifying that the bandwidth availability satisfies the batch task bandwidth usage threshold. Further, in one or more embodiments, the series of acts **2400** includes determining that the bandwidth availability indicated by the hardware usage metrics no longer satisfies the batch task bandwidth usage threshold and pausing the batch task based on determining that the bandwidth availability no longer satisfies the batch task bandwidth usage threshold.

(289) Moreover, in one or more embodiments, the series of acts includes monitoring hardware usage metrics for a hardware environment hosting a machine-learning model, extracting, at a content management system, workload data requesting execution of a training task using the machine-learning model, determining, at the content management system, that a training dataset for the training task is compatible with the hardware environment, and, based on determining that the training dataset is compatible with the hardware environment and the workload data, initiate the training task at the hardware environment based on a bandwidth availability indicated by the hardware usage metrics. Also, in one or more embodiments, the series of acts **2400** includes determining that the training dataset is compatible with the hardware environment by identifying that the training dataset is local to the hardware environment.

(290) Further, in one or more embodiments, the series of acts **2400** includes initiating the training task by determining that the training dataset is compatible with the hardware environment based on determining a data affinity for the hardware environment and the training dataset and initiating the training task based on determining the data affinity for the hardware environment and the training dataset. Moreover, in one or more embodiments, the series of acts **2400** includes forecasting a minimum-use time period where the hardware usage metrics indicate that an average bandwidth availability satisfies a training task bandwidth usage threshold, scheduling the training task for the minimum-use time period, and initiating the training task based on detecting that an actual bandwidth at the minimum-use time period satisfies the training task bandwidth usage threshold.

(291) In addition, in one or more embodiments, the series of acts **2400** includes pausing the training task based on detecting a change in the bandwidth availability indicated by the hardware usage metrics and resuming the training task based on detecting an additional change in the bandwidth availability indicated by the hardware usage metrics.

(292) As mentioned above, in certain embodiments, the intelligent selection and execution platform **102** trains or tunes a smart pocket machine-learning model (and/or a model selection machine-learning model or hardware allocation machine-learning model). In particular, in embodiments where the smart pocket machine-learning model is a series of gradient-boosted trees or a neural network, the intelligent selection and execution platform **102** utilizes an iterative training process to fit a smart machine-learning model by adjusting or adding decision trees or learning parameters that result in accurate model outputs. FIG. 25 illustrates training a smart pocket machine-learning model in accordance with one or more embodiments. Though only the smart pocket machine-learning model is depicted in FIG. 25, it is understood that the iterative training process may also be performed for any of the models described above (e.g., the model selection machine-learning model and/or the hardware allocation machine-learning model).

(293) As illustrated in FIG. 25, the intelligent selection and execution platform **102** accesses training task **2514**. Training task **2514** constitutes a task and/or workload data requesting the execution of a task and is used to train the smart pocket machine-learning model **2504**. The training task **2514** has a corresponding ground truth **2512** associated with it, where the ground truth **2512** indicates whether the training task **2514** was previously determined to be either a correct model selection (or hardware allocation and/or model training scheduling). For example, training task

2514 could be a selection of a machine-learning model (and, in some cases, a fallback machine-learning model) that a team of researchers determined or confirmed as either correct or incorrect. As another example, training task **2514** could be a hardware allocation that a team of researchers determined or confirmed as either correct or incorrect. Accordingly, in some cases, the intelligent selection and execution platform **102** treats the ground truth **2512** as a ground truth for training the smart pocket machine-learning model **2504**.

(294) As further illustrated in FIG. 25, the intelligent selection and execution platform **102** provides training features **2502** associated with the training task **2514** to the smart pocket machine-learning model **2504** and utilizes the smart pocket machine-learning model **2504** to generate a training model output **2508** based on the training features **2502**. As the name indicates, the training features **2502** represent features associated with the training task **2514** that are used for training the smart pocket machine-learning model **2504**. Accordingly, the training features **2502** can constitute a feature used as input for the smart pocket machine-learning model **2504**. In some embodiments, the smart pocket machine-learning model **2504** generates training model output **2508**, including a model selection, hardware allocation, and/or model training scheduling for the training task **2514**. The training model output **2508** can accordingly take the form of any of the model outputs described above.

(295) As further illustrated in FIG. 25, the intelligent selection and execution platform **102** utilizes a loss function to compare the training model output **2508** and the ground truth **2512** (e.g., to determine an error or a measure of loss between them). For instance, in cases where the smart pocket machine-learning model **2504** is an ensemble of gradient-boosted trees, the intelligent selection and execution platform **102** utilizes a mean squared error loss function (e.g., for regression) and/or a logarithmic loss function (e.g., for classification) as the loss function **2510**. By contrast, in embodiments where the smart pocket machine-learning model **2504** is a neural network, the intelligent selection and execution platform can utilize a cross-entropy loss function, an L1 loss function, or a mean squared error loss function as the loss function **2510**. For example, the intelligent selection and execution platform **102** utilizes the loss function **2510** to determine a difference between the training model output **2508** and the ground truth **2512**.

(296) As further illustrated in FIG. 25, the intelligent selection and execution platform performs model fitting **2506**. In particular, the intelligent selection and execution platform **102** fits the smart pocket machine-learning model **2504** based on loss from the loss function **2510**. For instance, the intelligent selection and execution platform **102** performs modifications or adjustments to the smart pocket machine-learning model **2504** to reduce the measure of loss from the loss function **2510** for a subsequent training iteration.

(297) For gradient-boosted trees, for example, the intelligent selection and execution platform **102** trains the smart pocket machine-learning model **2504** on the gradients of errors determined by the loss function **2510**. For instance, the intelligent selection and execution platform solves a convex optimization problem (e.g., of infinite dimensions) while regularizing the objective to avoid overfitting. In certain implementations, the intelligent selection and execution platform **102** scales the gradients to emphasize corrections to under-represented classes (e.g., fraud classifications or non-fraud classifications).

(298) In some embodiments, the intelligent selection and execution platform **102** adds a new weak learner (e.g., a new boosted tree) to the smart pocket machine-learning model **2504** for each successive training iteration as part of solving the optimization problem. For example, the intelligent selection and execution platform **102** finds a feature that minimizes a loss from the loss function **2510** and either adds the feature to the current iteration's tree or starts to build a new tree with the feature

(299) In addition to, or in the alternative, gradient-boosted decision trees, the intelligent selection and execution platform **102** trains a logistic regression to learn parameters for generating one or more fraud predictions, such as a fraud score indicating a probability of fraud. To avoid overfitting,

the intelligent selection and execution platform **102** further regularizes based on hyperparameters such as the learning rate, stochastic gradient boosting, the number of trees, the tree depth (s), complexity penalization, and L1/L2 regularization.

(300) In embodiments where the smart pocket machine-learning model **2504** is a neural network, the intelligent selection and execution platform **102** performs the model fitting **2506** by modifying internal parameters (e.g., weights) of the smart pocket machine-learning model **2504** to reduce the measure of loss for the loss function **2510**. Indeed, the intelligent selection and execution platform **102** modifies how the smart pocket machine-learning model **2504** analyzes and passes data between layers and neurons by modifying the internal network parameters. Thus, over multiple iterations, the intelligent selection and execution platform **102** improves the accuracy of the smart pocket machine-learning model **2504**.

(301) Indeed, in some cases, the intelligent selection and execution platform **102** repeats the training process illustrated in FIG. **25** for multiple iterations. For example, the intelligent selection and execution platform **102** repeats the iterative training by selecting a new set of training features for each training digital claim along with a corresponding fraud action label. The intelligent selection and execution platform **102** further generates a new set of training fraud predictions for each iteration. As described above, the intelligent selection and execution platform **102** also compares a training fraud prediction at each iteration with the corresponding training action label and further performs model fitting **2506**. The intelligent selection and execution platform **102** repeats this process until the smart pocket machine-learning model **2504** generates training fraud predictions that result in fraud predictions that satisfy a threshold measure of loss.

(302) Embodiments of the present disclosure may comprise or utilize a special purpose or general-purpose computer including computer hardware, such as, for example, one or more processors and system memory, as discussed in greater detail below. Embodiments within the scope of the present disclosure also include physical and other computer-readable media for carrying or storing computer-executable instructions and/or data structures. In particular, one or more of the processes described herein may be implemented at least in part as instructions embodied in a non-transitory computer-readable medium and executable by one or more computing devices (e.g., any of the media content access devices described herein). In general, a processor (e.g., a microprocessor) receives instructions, from a non-transitory computer-readable medium, (e.g., memory), and executes those instructions, thereby performing one or more processes, including one or more of the processes described herein.

(303) Computer-readable media can be any available media that can be accessed by a general purpose or special purpose computer system. Computer-readable media that store computer-executable instructions are non-transitory computer-readable storage media (devices). Computer-readable media that carry computer-executable instructions are transmission media. Thus, by way of example, and not limitation, embodiments of the disclosure can comprise at least two distinctly different kinds of computer-readable media: non-transitory computer-readable storage media (devices) and transmission media.

(304) Non-transitory computer-readable storage media (devices) includes RAM, ROM, EEPROM, CD-ROM, solid state drives (“SSDs”) (e.g., based on RAM), Flash memory, phase-change memory (“PCM”), other types of memory, other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store desired program code means in the form of computer-executable instructions or data structures and which can be accessed by a general purpose or special purpose computer.

(305) A “network” is defined as one or more data links that enable the transport of electronic data between computer systems and/or modules and/or other electronic devices. When information is transferred or provided over a network or another communications connection (either hardwired, wireless, or a combination of hardwired or wireless) to a computer, the computer properly views the connection as a transmission medium. Transmissions media can include a network and/or data

links which can be used to carry desired program code means in the form of computer-executable instructions or data structures and which can be accessed by a general purpose or special purpose computer. Combinations of the above should also be included within the scope of computer-readable media.

(306) Further, upon reaching various computer system components, program code means in the form of computer-executable instructions or data structures can be transferred automatically from transmission media to non-transitory computer-readable storage media (devices) (or vice versa). For example, computer-executable instructions or data structures received over a network or data link can be buffered in RAM within a network interface module (e.g., a “NIC”), and then eventually transferred to computer system RAM and/or to less volatile computer storage media (devices) at a computer system. Thus, it should be understood that non-transitory computer-readable storage media (devices) can be included in computer system components that also (or even primarily) utilize transmission media.

(307) Computer-executable instructions comprise, for example, instructions and data which, when executed by a processor, cause a general-purpose computer, special purpose computer, or special purpose processing device to perform a certain function or group of functions. In some embodiments, computer-executable instructions are executed by a general-purpose computer to turn the general-purpose computer into a special purpose computer implementing elements of the disclosure. The computer-executable instructions may be, for example, binaries, intermediate format instructions such as assembly language, or even source code. Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the described features or acts described above. Rather, the described features and acts are disclosed as example forms of implementing the claims.

(308) Those skilled in the art will appreciate that the disclosure may be practiced in network computing environments with many types of computer system configurations, including, personal computers, desktop computers, laptop computers, message processors, hand-held devices, multi-processor systems, microprocessor-based or programmable consumer electronics, network PCs, minicomputers, mainframe computers, mobile telephones, PDAs, tablets, pagers, routers, switches, and the like. The disclosure may also be practiced in distributed system environments where local and remote computer systems, which are linked (either by hardwired data links, wireless data links, or by a combination of hardwired and wireless data links) through a network, both perform tasks. In a distributed system environment, program modules may be located in both local and remote memory storage devices.

(309) Embodiments of the present disclosure can also be implemented in cloud computing environments. As used herein, the term “cloud computing” refers to a model for enabling on-demand network access to a shared pool of configurable computing resources. For example, cloud computing can be employed in the marketplace to offer ubiquitous and convenient on-demand access to the shared pool of configurable computing resources. The shared pool of configurable computing resources can be rapidly provisioned via virtualization and released with low management effort or service provider interaction, and then scaled accordingly.

(310) A cloud-computing model can be composed of various characteristics such as, for example, on-demand self-service, broad network access, resource pooling, rapid elasticity, measured service, and so forth. A cloud-computing model can also expose various service models, such as, for example, Software as a Service (“SaaS”), Platform as a Service (“PaaS”), and Infrastructure as a Service (“IaaS”). A cloud-computing model can also be deployed using different deployment models such as private cloud, community cloud, public cloud, hybrid cloud, and so forth. In addition, as used herein, the term “cloud-computing environment” refers to an environment in which cloud computing is employed.

(311) FIG. 26 illustrates a block diagram of an example computing device 2600 that may be

configured to perform one or more of the processes described above. One will appreciate that one or more computing devices, such as the computing device **2600** may represent the computing devices described above (e.g., client device(s) **112**, server(s) **110**, third-party server(s) **118**, and third-party server(s) **124**). In one or more embodiments, the computing device **2600** may be a mobile device (e.g., a mobile telephone, a smartphone, a PDA, a tablet, a laptop, a camera, a tracker, a watch, a wearable device, etc.). In some embodiments, the computing device **2600** may be a non-mobile device (e.g., a desktop computer or another type of client device). Further, the computing device **2600** may be a server device that includes cloud-based processing and storage capabilities.

(312) As shown in FIG. **26**, the computing device **2600** can include one or more processor(s) **2602**, memory **2604**, a storage device **2606**, input/output interfaces **2608** (or “I/O interfaces **2608**”), and a communication interface **2610**, which may be communicatively coupled by way of a communication infrastructure (e.g., bus **2612**). While the computing device **2600** is shown in FIG. **26**, the components illustrated in FIG. **26** are not intended to be limiting. Additional or alternative components may be used in other embodiments. Furthermore, in certain embodiments, the computing device **2600** includes fewer components than those shown in FIG. **26**. Components of the computing device **2600** shown in FIG. **26** will now be described in additional detail.

(313) In particular embodiments, the processor(s) **2602** includes hardware for executing instructions, such as those making up a computer program. As an example, and not by way of limitation, to execute instructions, the processor(s) **2602** may retrieve (or fetch) the instructions from an internal register, an internal cache, memory **2604**, or a storage device **2606** and decode and execute them.

(314) The computing device **2600** includes memory **2604**, which is coupled to the processor(s) **2602**. The memory **2604** may be used for storing data, metadata, and programs for execution by the processor(s). The memory **2604** may include one or more of volatile and non-volatile memories, such as Random-Access Memory (“RAM”), Read-Only Memory (“ROM”), a solid-state disk (“SSD”), Flash, Phase Change Memory (“PCM”), or other types of data storage. The memory **2604** may be internal or distributed memory.

(315) The computing device **2600** includes a storage device **2606** includes storage for storing data or instructions. As an example, and not by way of limitation, the storage device **2606** can include a non-transitory storage medium described above. The storage device **2606** may include a hard disk drive (HDD), flash memory, a Universal Serial Bus (USB) drive or a combination these or other storage devices.

(316) As shown, the computing device **2600** includes one or more I/O interfaces **2608**, which are provided to allow a user to provide input to (such as user strokes), receive output from, and otherwise transfer data to and from the computing device **2600**. These I/O interfaces **2608** may include a mouse, keypad or a keyboard, a touch screen, camera, optical scanner, network interface, modem, other known I/O devices or a combination of such I/O interfaces **2608**. The touch screen may be activated with a stylus or a finger.

(317) The I/O interfaces **2608** may include one or more devices for presenting output to a user, including, but not limited to, a graphics engine, a display (e.g., a display screen), one or more output drivers (e.g., display drivers), one or more audio speakers, and one or more audio drivers. In certain embodiments, I/O interfaces **2608** are configured to provide graphical data to a display for presentation to a user. The graphical data may be representative of one or more graphical user interfaces and/or any other graphical content as may serve a particular implementation.

(318) The computing device **2600** can further include a communication interface **2610**. The communication interface **2610** can include hardware, software, or both. The communication interface **2610** provides one or more interfaces for communication (such as, for example, packet-based communication) between the computing device and one or more other computing devices or one or more networks. As an example, and not by way of limitation, communication interface **2610**

may include a network interface controller (NIC) or network adapter for communicating with an Ethernet or other wire-based network or a wireless NIC (WNIC) or wireless adapter for communicating with a wireless network, such as a WI-FI. The computing device **2600** can further include a bus **2612**. The bus **2612** can include hardware, software, or both that connects components of computing device **2600** to each other.

(319) FIG. **27** is a schematic diagram illustrating environment **2700** within which one or more implementations of the intelligent selection and execution platform **102** can be implemented. For example, the intelligent selection and execution platform **102** may be part of a content management system **2702** (e.g., the content management system **108**). Content management system **2702** may generate, store, manage, receive, and send digital content (such as digital content items). For example, content management system **2702** may send and receive digital content to and from client devices **2706** by way of network **2704**. In particular, content management system **2702** can store and manage a collection of digital content. Content management system **2702** can manage the sharing of digital content between computing devices associated with a plurality of users. For instance, content management system **2702** can facilitate a user sharing a digital content with another user of content management system **2702**.

(320) In particular, content management system **2702** can manage synchronizing digital content across multiple client devices **2706** associated with one or more users. For example, a user may edit digital content using client device **2706**. The content management system **2702** can cause client device **2706** to send the edited digital content to content management system **2702**. Content management system **2702** then synchronizes the edited digital content on one or more additional computing devices.

(321) In addition to synchronizing digital content across multiple devices, one or more implementations of content management system **2702** can provide an efficient storage option for users that have large collections of digital content. For example, content management system **2702** can store a collection of digital content on content management system **2702**, while the client device **2706** only stores reduced-sized versions of the digital content. A user can navigate and browse the reduced-sized versions (e.g., a thumbnail of a digital image) of the digital content on client device **2706**. In particular, one way in which a user can experience digital content is to browse the reduced-sized versions of the digital content on client device **2706**.

(322) Another way in which a user can experience digital content is to select a reduced-size version of digital content to request the full- or high-resolution version of digital content from content management system **2702**. In particular, upon a user selecting a reduced-sized version of digital content, client device **2706** sends a request to content management system **2702** requesting the digital content associated with the reduced-sized version of the digital content. Content management system **2702** can respond to the request by sending the digital content to client device **2706**. Client device **2706**, upon receiving the digital content, can then present the digital content to the user. In this way, a user can have access to large collections of digital content while minimizing the amount of resources used on client device **2706**.

(323) Client device **2706** may be a desktop computer, a laptop computer, a tablet computer, a personal digital assistant (PDA), an in- or out-of-car navigation system, a handheld device, a smart phone or other cellular or mobile phone, or a mobile gaming device, other mobile device, or other suitable computing devices. Client device **2706** may execute one or more client applications, such as a web browser (e.g., Microsoft Windows Internet Explorer, Mozilla Firefox, Apple Safari, Google Chrome, Opera, etc.) or a native or special-purpose client application (e.g., Dropbox Paper for iPhone or iPad, Dropbox Paper for Android, etc.), to access and view content over network **2704**.

(324) Network **2704** may represent a network or collection of networks (such as the Internet, a corporate intranet, a virtual private network (VPN), a local area network (LAN), a wireless local area network (WLAN), a cellular network, a wide area network (WAN), a metropolitan area

network (MAN), or a combination of two or more such networks) over which client devices may access content management system 2702.

(325) In the foregoing specification, the invention has been described with reference to specific example embodiments thereof. Various embodiments and aspects of the invention(s) are described with reference to details discussed herein, and the accompanying drawings illustrate the various embodiments. The description above and drawings are illustrative of the invention and are not to be construed as limiting the invention. Numerous specific details are described to provide a thorough understanding of various embodiments of the present invention.

(326) The present invention may be embodied in other specific forms without departing from its spirit or essential characteristics. The described embodiments are to be considered in all respects only as illustrative and not restrictive. For example, the methods described herein may be performed with less or more steps/acts or the steps/acts may be performed in differing orders. Additionally, the steps/acts described herein may be repeated or performed in parallel to one another or in parallel to different instances of the same or similar steps/acts. The scope of the invention is, therefore, indicated by the appended claims rather than by the foregoing description. All changes that come within the meaning and range of equivalency of the claims are to be embraced within their scope.

Claims

1. A computer-implemented method comprising: receiving, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model; extracting, from the workload data, workload features defining characteristics of the task; determining task routing metrics for a plurality of hardware environments hosted in respective network environments and a plurality of machine-learning models in respective network environments; determining a historical quality metric indicating how one or more machine learning models of the plurality of machine-learning models will execute the task based on one or more user feedback metrics; and selecting, based on an output of a model selection machine-learning model that utilizes the historical quality metric, from the plurality of hardware environments and from the plurality of machine-learning models, a designated hardware environment and a designated machine-learning model for executing the task based on the workload features and the task routing metrics.

2. The computer-implemented method of claim 1, wherein determining task routing metrics for the plurality of hardware environments further comprises: determining a hardware state for each hardware environment of the plurality of hardware environments; and wherein selecting, from the plurality of hardware environments, the designated hardware environment for executing the task is based in part on the hardware state of the designated hardware environment.

3. The computer-implemented method of claim 1, wherein determining task routing metrics for the plurality of hardware environments further comprises: determining, based on the workload features, a financial cost metric, an execution time metric, an execution cost metric, or a model fit metric for executing the task on each hardware environment of the plurality of hardware environments; and wherein selecting, from the plurality of hardware environments, the designated hardware environment for executing the task is based on two or more of the financial cost metric, the execution time metric, the execution cost metric, and the model fit metric.

4. The computer-implemented method of claim 1, wherein selecting the designated hardware environment further comprises: analyzing the workload data and the task routing metrics for the plurality of hardware environments to determine an optimal machine-learning model for executing the task from the plurality of machine-learning models; and selecting the designated machine-learning model for executing the task based on determining that the designated machine-learning model is the optimal machine-learning model from the plurality of machine-learning models.

5. The computer-implemented method of claim 4, wherein determining that the designated machine-learning model is the optimal machine-learning model further comprises: generating, utilizing a smart pocket machine-learning model, an optimization metric; and determining that the designated machine-learning model is the optimal machine-learning model based on the optimization metric.
6. The computer-implemented method of claim 1, wherein extracting workload features defining characteristics of the task further comprises determining an estimated processing requirement and an estimated storage requirement for executing the task.
7. The computer-implemented method of claim 1, wherein selecting the designated hardware environment further comprises: analyzing the workload data and the task routing metrics for the plurality of hardware environments of the plurality of hardware environments, wherein the plurality of hardware environments comprises one or more hardware environments and one or more third-party hardware environments; and selecting the designated hardware environment for executing the task based on analyzing the workload data and the task routing metrics.
8. The computer-implemented method of claim 1, wherein selecting the designated hardware environment further comprises: performing probabilistic load balancing of each hardware environments of the plurality of hardware environments for executing the task; and selecting the designated hardware environment based on the probabilistic load balancing.
9. The computer-implemented method of claim 1, wherein selecting the designated hardware environment further comprises: accessing the one or more user feedback metrics about executing tasks using one or more machine-learning models of the plurality of machine-learning models; determining the historical quality metric indicating how the one or more machine-learning models of the plurality of machine-learning models will execute the task based on the one or more user feedback metrics; and selecting the designated machine-learning model for executing the task based on the historical quality metric.
10. The computer-implemented method of claim 1, wherein selecting the designated hardware environment further comprises: identifying a capability or specialty associated with one or more machine-learning models of the plurality of machine-learning models; and selecting the designated machine-learning model for executing the task based on alignment of the capability or the specialty of the designated hardware environment with the task.
11. A non-transitory computer-readable medium storing instructions that, when executed by at least one processor, cause a computer system to: receive, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model; extract, from the workload data, workload features defining characteristics of the task; determine task routing metrics for a plurality of hardware environments hosted in respective network environments and a plurality of machine-learning models in respective network environments; determine a historical quality metric indicating how one or more machine learning models of the plurality of machine-learning models will execute the task based on one or more user feedback metrics; and select, based on an output of a model selection machine-learning model that utilizes the historical quality metric, from the plurality of hardware environments and from the plurality of machine-learning models and based on the workload features and the task routing metrics: a first designated hardware environment and a first designated machine-learning model for executing a first portion of the task; and a second designated hardware environment and a second designated machine-learning model for executing a second portion of the task.
12. The non-transitory computer-readable medium of claim 11, further comprising instructions that, when executed by the at least one processor, cause the computer system to: determine a hardware state for each hardware environment of the plurality of hardware environments; determine a model state for each machine-learning model for the plurality of machine-learning models; select the first designated hardware environment for executing the first portion of the task is based on a first hardware state for the first designated hardware environment and a first model state for a first

machine-learning model; and select the second designated hardware environment for executing the second portion of the task is based on a second hardware state for the second designated hardware environment and a second model state for a second machine-learning model.

13. The non-transitory computer-readable medium of claim 11, further comprising instructions that, when executed by the at least one processor, cause the computer system to: determine, based on the workload features, a financial cost metric, an execution time metric, an execution cost metric, or a model fit metric for executing the task on each machine-learning model of the plurality of machine-learning models; select the first designated machine-learning model for executing the first portion of the task based on two or more of the financial cost metric, the execution time metric, the execution cost metric, or the model fit metric for executing the first portion of the task on the first designated machine-learning model; and select the second designated machine-learning model for executing the second portion of the task based on two or more of the financial cost metric, the execution time metric, the execution cost metric, or the model fit metric for executing the second portion of the task on the second designated machine-learning model.

14. The non-transitory computer-readable medium of claim 11, further comprising instructions that, when executed by the at least one processor, cause the computer system to: analyze the workload data and the task routing metrics for the plurality of hardware environments; select the first designated hardware environment for executing the first portion of the task based on determining that the first designated hardware environment is a first optimal hardware environment from the plurality of hardware environments; and select the second designated hardware environment for executing the second portion of the task based on determining that the second designated hardware environment is a second optimal hardware environment for executing the second portion of the task.

15. The non-transitory computer-readable medium of claim 11, further comprising instructions that, when executed by the at least one processor, cause the computer system to: perform probabilistic load balancing of each hardware environment of the plurality of hardware environments for executing the task; and select the first designated hardware environment for executing the first portion of the task and the second designated hardware environment for executing the second portion of the task based on the probabilistic load balancing.

16. A system comprising: at least one processor; and at least one non-transitory computer-readable storage medium storing instructions that, when executed by the at least one processor, cause the system to: receive, from a device connected by a network to a content management system, workload data requesting execution of a task using a machine-learning model; extract, from the workload data, workload features defining characteristics of the task; determine task routing metrics for a plurality of hardware environments hosted in respective network environments and a plurality of machine-learning models in respective network environments; determine a historical quality metric indicating how one or more machine learning models of the plurality of machine-learning models will execute the task based on a user feedback metric; and select, utilizing an output of a smart pocket machine-learning model that utilizes the historical quality metric, a designated hardware environment from the plurality of hardware environments and a designated machine-learning model from the plurality of machine-learning models for executing the task based on the workload features and the task routing metrics.

17. The system of claim 16, further comprising instructions that, when executed by the at least one processor, cause the system to: utilize the smart pocket machine-learning model to compare each hardware environment of the plurality of hardware environments and each machine learning model of the plurality of machine-learning models based on the workload features and the task routing metrics; and select the designated hardware environment from the plurality of hardware environments and the designated machine-learning model from the plurality of machine-learning models based on an output of the smart pocket machine-learning model.

18. The system of claim 16, further comprising instructions that, when executed by the at least one

processor, cause the system to: receive additional task routing metrics for an additional hardware environment; and update parameters of the smart pocket machine-learning model based on the additional task routing metrics of the additional hardware environment.

19. The system of claim 16, further comprising instructions that, when executed by the at least one processor, cause the system to select the designated hardware environment by: generate, utilizing the smart pocket machine-learning model, an optimization metric for each hardware environment of the plurality of hardware environments; and select the designated hardware environment based on the optimization metric.

20. The system of claim 16, further comprising instructions that, when executed by the at least one processor, cause the system to: receive user feedback data indicating a user satisfaction with execution of the task; and update parameters of the smart pocket machine-learning model based on the user feedback data.
