

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386669
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Hasabnis; Niranjana et al.

Methods and apparatus to determine execution cost

Abstract

Methods, apparatus, systems, and articles of manufacture to determine execution cost are disclosed. An example apparatus includes memory; instructions included in the apparatus; and processor circuitry to execute the instruction to: cause a plurality of instructions corresponding to a mnemonic to be executed; determine an average execution cost of the plurality of instructions; determine a standard deviation of execution costs of the plurality of instructions; and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

Inventors: Hasabnis; Niranjana (Fremont, CA), Gottschlich; Justin (Santa Clara, CA), Tithi; Jesmin Jahan (San Jose, CA), Venkat; Anand (San Jose, CA)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 80740391

Assignee: Intel Corporation (Santa Clara, CA)

Appl. No.: 17/541016

Filed: December 02, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220091895 A1	Mar. 24, 2022

Publication Classification

Int. Cl.: G06F9/50 (20060101); G06F9/30 (20180101); G06F9/38 (20180101); G06F9/455 (20180101)

U.S. Cl.:

CPC **G06F9/5044** (20130101); **G06F9/30101** (20130101); **G06F9/3836** (20130101);
G06F9/4552 (20130101); **G06F9/5016** (20130101); **G06F9/5038** (20130101);

Field of Classification Search

CPC: G06F (8/443); G06F (9/4552); G06F (9/5044); G06F (9/5016); G06F (9/5038); G06F (9/30101); G06F (9/3836)

USPC: 717/151-161

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9811335	12/2016	Khainovski	N/A	G06F 9/30196
2011/0231830	12/2010	Udayakumaran	717/160	G06F 8/4441
2022/0155769	12/2021	Zhou	N/A	H04M 15/00

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
104424363	12/2018	CN	G06F 16/20
106874332	12/2019	CN	G06F 16/2433

OTHER PUBLICATIONS

LLVM, “ILVM-MCA—LLVM Machine Code Analyzer,” LLVM Compiler Infrastructure, Synopsis, 2021, retrieved Jul. 5, 2023. [Online]. Available: <https://llvm.org/docs/CommandGuide/llvm-mca.html>, 18 pages. cited by applicant

Park, “Predicting program execution times by analyzing static and dynamic program paths,” Real Time Systems, vol. 5, Mar. 1993, 32 pages. cited by applicant

Fahringer et al., “A static parameter based performance prediction tool for parallel programs,” in Proceedings of the 7th international conference on Supercomputing, Jul. 19, 1993, 13 pages. cited by applicant

Rugina et al., “Predicting the running times of parallel programs by simulation,” in IEEE, Mar. 30, 1998, 7 pages. cited by applicant

Derumigny et al., “From micro-OPs to abstract resources: constructing a simpler CPU performance model through microbenchmarking,” arXiv, dated Sep. 16, 2021, 16 pages. cited by applicant

Sanchez et al., “ZSim: Fast and accurate microarchitectural simulation of thousand-core systems,” ACM SIGARCH Computer architecture news, Jun. 2013, 12 pages. cited by applicant

Xu et al., “Explainable AI: A Brief Survey on History, Research Areas, Approaches and Challenges,” Natural Language Processing and Chinese Computing: 8th CCF International Conference, NLPCC 2019, Oct. 9, 2019, 12 pages. cited by applicant

Li et al., “LightRNN: Memory and Computation-Efficient Recurrent Neural Networks,” 30th Conference on Neural Information Processing Systems (NIPS 2016), Barcelona, Spain, Dec. 5, 2016, 9 pages. cited by applicant

Gottschlich et al., “The three pillars of machine programming,” in 2nd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages, Philadelphia, Jun. 18, 2018, 12 pages. cited by applicant

Laukemann et al., "Automated instruction stream throughput prediction for Intel and AMD microarchitectures," in IEEE/ACM Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS), Nov. 12, 2018, 11 pages. cited by applicant

Intel, "Intel architecture code analyzer (IACA)," [Online]. Available: <https://software.intel.com/content/www/us/en/develop/articles/intel-architecture-code-analyzer.html>, 2019, 4 pages. cited by applicant

Abel et al., "Characterizing Latency, Throughput, and Port Usage of Instructions on Intel Microarchitectures," in Twenty-Fourth International Conference on Architectural Support for Programming Languages and Operating Systems, Apr. 13, 2019, 14 pages. cited by applicant

Mendis et al., "Ithemal: Accurate, Portable and Fast Basic Block Throughput Estimation using Deep Neural Networks," 36th International Conference on Machine Learning (ICML) 2019, May 30, 2019, 14 pages. cited by applicant

Blanco et al., "Predicting the performance of parallel programs," ELSEVIER, Parallel Computing, accepted Nov. 20, 2003, 20 pages. cited by applicant

Dauphin et al., PEPP: Performance Evaluation of Parallel Programs, Erlangen, User's Guide version 3.3, Sep. 1993, 78 pages. cited by applicant

Primary Examiner: Vo; Ted T.

Attorney, Agent or Firm: Hanley, Flight & Zimmerman, LLC.

Background/Summary

FIELD OF THE DISCLOSURE

(1) This disclosure relates generally to computing systems and, more particularly, to methods and apparatus to determine execution cost.

BACKGROUND

(2) Processor resources are needed to execute instructions to perform one or more tasks. The amount of processor resources needed to execute the instructions corresponds to an execution cost of the instructions. The execution cost may correspond to a total number of cycles and/or seconds to execute the program. Because different computer architectures are structured differently, a program designed for one type of computer architecture may have a different execution cost for another computer architecture. A computing system may use execution cost of code to attempt to optimize instruction, improve instructions, and/or otherwise lower the execution cost of the instructions. For example, machine programming may rely on execution cost when attempting to replace low-performance instructions with high-performance instructions based on the execution cost of both sets of instructions.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is an example environment to determine the execution cost in conjunction with examples disclosed herein.
- (2) FIG. 2 is a block diagram of an implementation of the execution cost mapping table determination circuitry and the basic block execution cost determination circuitry of FIG. 1.
- (3) FIGS. 3A and 3B illustrate a flowchart representative of example machine readable instructions that may be executed by example processor circuitry to implement the execution cost mapping

table determination circuitry of FIG. 2.

(4) FIG. 4 illustrates a flowchart representative of example machine readable instructions that may be executed by example processor circuitry to implement the basic block execution cost determination circuitry of FIG. 2.

(5) FIG. 5 is a block diagram of an example processing platform including processor circuitry structured to execute the example machine readable instructions of FIGS. 3A, 3B and/or 4 to implement the example computing device of FIG. 1.

(6) FIG. 6 is a block diagram of an example implementation of the processor circuitry of FIG. 5.

(7) FIG. 7 is a block diagram of another example implementation of the processor circuitry of FIG. 5.

(8) FIG. 8 is a block diagram of an example software distribution platform (e.g., one or more servers) to distribute software (e.g., software corresponding to the example machine readable instructions of FIGS. 3A, 3B and/or 4) to client devices associated with end users and/or consumers (e.g., for license, sale, and/or use), retailers (e.g., for sale, re-sale, license, and/or sub-license), and/or original equipment manufacturers (OEMs) (e.g., for inclusion in products to be distributed to, for example, retailers and/or to other end users such as direct buy customers).

DETAILED DESCRIPTION

(9) The figures are not to scale. Instead, the thickness of the layers or regions may be enlarged in the drawings. As used herein, connection references (e.g., attached, coupled, connected, and joined) may include intermediate members between the elements referenced by the connection reference and/or relative movement between those elements unless otherwise indicated. As such, connection references do not necessarily infer that two elements are directly connected and/or in fixed relation to each other. As used herein, stating that any part is in “contact” with another part is defined to mean that there is no intermediate part between the two parts.

(10) Unless specifically stated otherwise, descriptors such as “first,” “second,” “third,” etc., are used herein without imputing or otherwise indicating any meaning of priority, physical order, arrangement in a list, and/or ordering in any way, but are merely used as labels and/or arbitrary names to distinguish elements for ease of understanding the disclosed examples. In some examples, the descriptor “first” may be used to refer to an element in the detailed description, while the same element may be referred to in a claim with a different descriptor such as “second” or “third.” In such instances, it should be understood that such descriptors are used merely for identifying those elements distinctly that might, for example, otherwise share a same name. As used herein, “approximately” and “about” refer to dimensions that may not be exact due to manufacturing tolerances and/or other real world imperfections. As used herein “substantially real time” refers to occurrence in a near instantaneous manner recognizing there may be real world delays for computing time, transmission, etc. Thus, unless otherwise specified, “substantially real time” refers to real time ± 1 second. As used herein, the phrase “in communication,” including variations thereof, encompasses direct communication and/or indirect communication through one or more intermediary components, and does not require direct physical (e.g., wired) communication and/or constant communication, but rather additionally includes selective communication at periodic intervals, scheduled intervals, aperiodic intervals, and/or one-time events.

(11) As used herein, “processor circuitry” is defined to include (i) one or more special purpose electrical circuits structured to perform specific operation(s) and including one or more semiconductor-based logic devices (e.g., electrical hardware implemented by one or more transistors), and/or (ii) one or more general purpose semiconductor-based electrical circuits programmed with instructions to perform specific operations and including one or more semiconductor-based logic devices (e.g., electrical hardware implemented by one or more transistors). Examples of processor circuitry include programmed microprocessors, Field Programmable Gate Arrays (FPGAs) that may instantiate instructions, Central Processor Units (CPUs), Graphics Processor Units (GPUs), Digital Signal Processors (DSPs), XPU, or

microcontrollers and integrated circuits such as Application Specific Integrated Circuits (ASICs). For example, an XPU may be implemented by a heterogeneous computing system including multiple types of processor circuitry (e.g., one or more FPGAs, one or more CPUs, one or more GPUs, one or more DSPs, etc., and/or a combination thereof) and application programming interface(s) (API(s)) that may assign computing task(s) to whichever one(s) of the multiple types of the processing circuitry is/are best suited to execute the computing task(s).

(12) Accurate and efficient determination of a program's cost (e.g., total number of clock cycles and/or seconds to complete a program) is desirable in many different applications (e.g., code optimization, design of hardware (e.g., compilers), design of software (e.g., binary translators), code recommendation systems, schedulers on cloud or clusters, etc.). For example, compiler backends may select instructions by relying on an estimation of execution cost. To determine execution cost of instructions, some techniques simply execute them on the actual hardware. However, such techniques are expensive and time consuming. Other techniques have been used to estimate execution cost using models and/or other tools. However, such techniques may be inaccurate and/or require large maintenance overhead. Some techniques use artificial intelligence (AI) to predict the execution cost of a set of instructions. Although such techniques are accurate, such techniques are computationally-intensive and, moreover, AI-based techniques do not provide an explanation for a result. Accordingly, the reason for why the AI-based techniques determined that the code corresponds to an execution cost is unknown.

(13) Examples disclosed herein develop an approach that is accurate, explainable, and executes using less processor resources than the above-mentioned techniques. Examples disclosed herein include a server that develops architecture-specific mapping tables that map converted instructions (e.g., abstractions of instructions based on the leading mnemonic and/or function of an instruction line) to an average execution cost and standard deviation cost. Using examples disclosed herein, the server can obtain programs and/or basic blocks of code from storage (e.g., open-source repositories, crowd-sourced repositories, closed-source repositories, local storage, external stored, local databases, external databases, etc.), execute lines of the code to identify execution costs corresponding to the function and/or mnemonic of the line of code and map the mnemonic to the average execution cost and the standard deviation of the execution cost. For example, examples disclosed herein may identify 10 instances of instructions that begin with the mnemonic “add,” execute the 10 instances on a CPU corresponding to a particular architecture and determine that the average cycles needed to execute the instruction was 3 with standard deviation of 1. In this manner, examples disclosed herein map the mnemonic “add” to a “3” average and a “1” standard deviation. As used herein, a basic block is a sequence of instructions or program, where the control flows sequentially through the sequence.

(14) After a threshold number of programs and/or blocks of code have been mapped, examples disclosed herein can deploy the mapping table to computing devices within a network that may or may not correspond to the particular computer architecture. In this manner, the computing system can use the mapping to estimate the execution cost of a program corresponding to a basic block by identifying mnemonic for each line of the program and determining the corresponding average and/or standard deviation of the execution cost using the mapping table, thereby requiring far less resources and time than the above-mentioned techniques. In some examples, the architecture of the computing device during deployment could be different than the architecture of the mapping table. Additionally, the mapping provides the reasoning and/or explainability for the cost of a program line-by-line.

(15) FIG. 1 is a block diagram of an example environment **100** described in conjunction with examples disclosed herein. The example environment **100** includes an example server **102**, example execution cost mapping table determination circuitry **104**, example architecture-specific CPU(s) **105**, an example network **106**, an example computing device **108**, an example mapping table storage **112**, an example CPU **114**, and an example basic block execution cost determination

circuitry **116**.

(16) The example server(s) **102** of FIG. **1** may be server(s) and/or any other computing device(s) that generates mapping tables. The example server **103** may access blocks of code stored locally or externally (e.g., a code repository/storage/database stored locally or externally). After blocks of code have been obtained, the example execution cost mapping table determination circuitry **104** processes the blocks of code to generate a mapping table that links average and standard deviation of execution cost to a mnemonic and/or function of a line of instruction. As further described below, the execution cost mapping table determination circuitry **104** processes the lines of the blocks of code and generates the mnemonic by abstracting out specific numerical and/or register values. In this manner, instructions corresponding to the same mnemonic are executed by one or more of the architecture-specific CPU(s) **105** and the execution cost is identified. In some examples, the architecture of the server **102** and the computing device **108** are different. For example, the computing device **108** may be a client device that has CPU for a first architecture but is optimizing a program written for a second architecture using the mapping table for the second architecture. This may occur because compilers may support a cross compilation mode, where a compiler running on the second architecture can generate or optimize a program written for the first architecture. The execution cost for each instruction corresponding to the same mnemonic are averaged and a standard deviation is calculated. The mnemonic is stored in a mapping table in conjunction with the corresponding average and standard deviation. In this manner, the computing device **108** can use the entry in the mapping table to estimate execution cost based on a program that includes a line of code that corresponds to the mnemonic.

(17) Because execution cost for the same instruction may be different for different architectures, the example server **102** of FIG. **1** includes the different architecture-specific CPU(s) **105** to execute code according to different specific architectures. The architecture-specific CPU(s) **105** executes code and outputs execution cost corresponding to the executed code. For example, the architecture-specific CPU(s) **105** acts as a driver program with necessary boilerplate code around concrete assembly instructions to execute individual instructions in a loop to reach a steady state (e.g., 100 iterations) and measure execution cost based on the iterations. In this manner, the execution cost mapping table determination circuitry **104** can develop architecture-specific mapping tables that can be deployed to computing devices. The example execution cost mapping table determination circuitry **104** is further described below in conjunction with FIG. **2**.

(18) The example network **106** of FIG. **1** is a system of interconnected systems exchanging data. The example network **106** may be implemented using any type of public or private network such as, but not limited to, the Internet, a telephone network, a local area network (LAN), a cable network, and/or a wireless network. To enable communication via the network **106**, the example server(s) **102** and/or the computing device **108** includes a communication interface that enables a connection to an Ethernet, a digital subscriber line (DSL), a telephone line, a coaxial cable, or any wireless connection, etc. In some examples, the server(s) **102** and the example computing device **108** are connected via the example network **106**.

(19) The example computing device **108** of FIG. **1** obtains determined mapping tables from the server **102** (e.g., via the network **106**) and uses the mapping table to estimate the execution costs of programs. The example computing device **108** may be a server, an edge-based device, a cloud-based device, an internet of thing (IoT) device, a computer, a mobile device, a tablet, and/or any other type of computing device. The example computing device **108** includes the example network interface **110** to obtain the determined mapping table corresponding to the architecture of the computing device **108** from the server **102** via the network **106**. An obtained mapping table is stored in the mapping table storage **112** of the example computing device **108**. Additionally, the network interface **110** may obtain mapping table updates and/or replacements. In such examples, the updates and/or replacements are used to update the mapping table stored in the example mapping table storage **112**.

(20) The example CPU **114** of FIG. **1** executes instructions (e.g., a program and/or code) to perform one or more tasks. In some examples, the CPU **114** may instruct the basic block execution cost determination circuitry **116** to determine the cost of program, block of code, and/or instruction. The basic block execution cost determination circuitry **116** of FIG. **1** determines the cost by traversing through the code line-by-line and determining a range of the execution cost of each line of code based on the execution cost and standard deviation corresponding to the mnemonic of the line using the mapping table stored in the example mapping table storage **112**. For example, when the basic block execution cost determination circuitry **116** processes a line of code, the cost determination circuitry determines the mnemonic and/or function part of the line of code (e.g., “add,” “div,” “store,” etc.). After the mnemonic is obtained, the basic block execution cost determination circuitry **116** accesses the mapping table in the mapping table storage **112** to identify an average and standard deviation for the determined mnemonic. The example basic block execution cost determination circuitry **116** determines an execution cost maximum and minimum based on the average and standard deviation. After the maximum and minimum execution costs of each line is determined, the example basic block execution cost determination circuitry **116** determines the average execution cost and standard deviation cost for the code based on the maximum and minimum execution costs. If the basic block execution cost determination circuitry **116** is unable to match the determined mnemonic with a mnemonic in the mapping table, the basic block execution cost determination circuitry **116** may determine the average and standard deviation of the line of code based on present values (e.g., the mean, median, or mode of the average execution cost and/or the mean, median, or mode of the standard deviation execution cost across the mnemonics of the mapping table). The example basic block execution cost determination circuitry **116** is further described below in conjunction with FIG. **2**.

(21) FIG. **2** is a block diagram of an example implementation of the example execution cost mapping table determination circuitry **104** and the example basic block execution cost determination circuitry **116** of FIG. **1**. The example execution cost mapping table determination circuitry **104** includes an example network interface **200**, example architecture-specific compiler(s) **202**, example instruction converter circuitry **204**, example pair storage **206**, example cost determination circuitry **208**, and an example component interface **210**. The example basic block execution cost determination circuitry **116** includes an example component interface **212**, example instruction conversion circuitry **214**, example cost estimation circuitry **216**, and example calculation circuitry **218**.

(22) The example network interface **200** of FIG. **2** transmits mapping tables to computing device(s) (e.g., the computing device **108**) via the network **106**. In some examples, the network interface **200** obtains instructions, code, programs, and/or basic blocks from one or more external repositories, storage, and/or databases.

(23) The example architecture-specific compiler(s) **202** compiles high level code into assembly code corresponding to a target architecture. The example architecture-specific compiler(s) **202** may be a single compiler to compile high level code into multiple architectures or may be multiple compilers (e.g., a first compiler for a first architecture, a second compiler for a second architecture, etc.). The example architecture-specific compiler(s) **202** may be GNU compiler collection (GCC), Intel® C++ Compiler (ICC), low level virtual machine (LLVM), etc. that compile(s) input programs using different permutations of compiler options (e.g., -O2, -march=T, etc.). Additionally, the example architecture-specific compiler(s) **202** may convert commands in makefiles (e.g., if available).

(24) The example instruction converter circuitry **204** of FIG. **2** converts lines of assembly code into mnemonic and operands. For example, the instruction converter circuitry **204** may convert “add \$2, %rax” into “add” and “\$2, %rax.” After converting in the mnemonic and operands, the example instruction convert circuitry **204** abstracts the command by removing the values and/or variables that correspond to numbers and/or registers and replaces them with tokens corresponding to the

value type. For example, operands that include numerical constants will be replaced with tokens corresponding to numerical type (e.g., integer, real number, Boolean, etc.) and operands that correspond to register will be replaced with a token corresponding to register. Additionally, the example instruction converter circuitry **204** pairs the assembly instruction with the corresponding abstraction and stores the pair in the example pair storage **206**. In some examples, before storing the pair in the example pair storage **206**, the example instruction converter circuitry **204** determines if the pair and/or the abstraction and token is already stored in the pair storage **206**, the instruction converter circuitry **204** may discard the pair (e.g., because a duplicate instruction may not add diversity to the stored pairs when determining execution cost information). The example instruction converter circuitry **204** may continue to process additional input programs until a threshold number of pairs are stored in the example pair storage **206**. The threshold may be based on user and/or manufacturer preferences.

(25) The example cost determination circuitry **208** of FIG. 2 accesses the example pair storage **206** to access one or more stored pairs. The example cost determination circuitry **208** may identify the pairs that correspond to the same mnemonic (e.g., with different operands) based on the abstraction. The cost determination circuitry **208** transmits (e.g., via the component interface **210**) the pairs of corresponding to the same mnemonic to the architecture-specific CPU **105** (e.g., corresponding a specific architecture that corresponds to the architecture of the architecture-specific compiler **202**) and obtains (e.g., via the component interface **210**) the execution cost for each of the instructions. The example cost determination circuitry **208** determines the average and standard deviation of the execution costs of the plurality of instructions corresponding to the same mnemonic. After the average and standard deviation corresponding to the mnemonic are determined, the cost determination circuitry **208** generates an entry for a mapping table to include the mnemonic and corresponding average and standard deviation cost.

(26) The example component interface **210** of FIG. 2 transmits instructions to be executed to the example architecture-specific CPU(s) **105**. Additionally, the example component interface **210** obtains execution cost of executing the instructions from the architecture-specific CPU(s) **105**. In some examples, the component interface **210** may obtain high level instructions, code, and/or basic blocks from internal storage of the server **102** (e.g., when the server **102** includes storage, one or more databases, and/or one or more repositories including code).

(27) The example component interface **212** of FIG. 2 of the example basic block execution cost determination circuitry **116** accesses mapping tables from the example mapping table storage **112**. As described above, the network interface **110** of the computing device **108** obtains the mapping table from the server **102** and stores the mapping table into the example mapping table storage **112**. Additionally, the example component interface **212** may obtain instructions to determine the execution cost of a program from the example CPU **114**. The example component interface **212** obtains a basic block, program, and/or code in assembly language to determine the execution cost. After the execution cost is determined, the example component interface **212** transmits the execution cost to the example CPU **114** or another device.

(28) The example instruction conversion circuitry **214** of FIG. 2 converts obtained and/or determined assembly language code into individual lines. In some examples, the CPU **114** includes a compiler to convert high level language into the assembly language. After the code is broken up into individual instructions, the example instruction conversion circuitry **214** identifies the mnemonic (e.g., function) of the individual instructions. For example, the instruction conversion circuitry **214** may identify the mnemonic “div” from the instructions “div \$2.”

(29) The example cost estimation circuitry **216** processes the mnemonics of the basic block, code, and/or program to determine an execution cost for the basic block, code, and/or program and/or a standard deviation for the basic block, code, and/or program. For example, for an instruction that corresponds to a “mov” mnemonic, the cost estimation circuitry **216** accesses the mapping table to identify an entry for the mnemonic “mov.” After the corresponding entry is identified, the cost

estimation circuitry **216** determines corresponding execution cost average (e.g., also referred to as range average, R_{avg}) and execution cost standard deviation (e.g., also referred to as range standard deviation, R_{stdev}) of the mnemonic from the mapping table. If the mnemonic of the instruction is not included in the mapping table, the example cost estimation circuitry **216** may use a preset value for the execution cost average and the execution standard deviation. In some examples, the cost estimation circuitry **216** selects the preset value based on the average execution cost and average standard deviation across the mnemonics from the mapping table.

(30) The example calculation circuitry **218** of FIG. 2 calculates the execution cost of the program (e.g., also referred to as a range summation) and a standard deviation for the execution cost of the program using the average execution cost and standard deviation of the execution cost of the instructions of the program, code, and/or basic block. For example, the calculation circuitry **218** may determine a minimum execution cost and/or cost range for a particular instruction using the below Equation 1, a maximum execution cost and/or cost range for the particular instruction using the below Equation 2, and the execution cost for the program (e.g., range summation) using the below Equation 3 (e.g., the sum of the average costs for the instructions of the program, code, and/or basic block).

$$\min(R) = R_{sub.avg} - R_{sub.stdev} \quad (\text{Equation 1})$$

$$\max(R) = R_{sub.avg} + R_{sub.stdev} \quad (\text{Equation 2})$$

$$\text{range_summation}(R_{sub.1}, \dots, R_{sub.n}) = R_{sub.1} + R_{sub.2} + \dots + R_{sub.n} \quad (\text{Equation 3})$$

(31) In the above Equations 1-3, R_{avg} is the average execution cost for an instruction and R_{stdev} is the standard deviation for the execution cost. Because standard deviation may provide valuable information that is not included in the range summation, the example calculation circuitry **218** also calculates the standard deviation of the minimum execution costs and maximum execution costs of the instructions of the program, code, and/or basic block, as shown in the below Equation 4.

$$\text{standard_deviation}(R_{sub.1}, \dots, R_{sub.n}) = \sigma(\min(R_{sub.1}), \dots, \min(R_{sub.n}), \max(R_{sub.1}), \dots, \max(R_{sub.n})) \quad (\text{Equation 4})$$

(32) The example calculation circuitry **218** transmits the execution cost and standard deviation of the program, code, and/or basic block to the CPU **114** and/or other component/device using the example component interface **212**.

(33) In some examples, the computational cost mapping table determination circuitry **104** includes means for identifying, means for generating, and means for converting. For example, the means for identify may be implemented by the instruction converter circuitry **204**, the means for generating may be implemented by the cost determination circuitry **208**, and the means for converting may be implemented by the instruction converter circuitry **204** or the architecture-specific compiler(s) **202**. In some examples, the architecture-specific compiler(s) **202**, the instruction converter circuitry **204**, and/or the cost determination circuitry **208** may be instantiated by processor circuitry such as the example processor circuitry **512** of FIG. 5. For instance, the architecture-specific compiler(s) **202**, the instruction converter circuitry **204**, and/or the cost determination circuitry **208** may be instantiated by the example general purpose processor circuitry **600** of FIG. 6 executing machine executable instructions such as that implemented by at least blocks of FIGS. 3A and 3B. In some examples, the architecture-specific compiler(s) **202**, the instruction converter circuitry **204**, and/or the cost determination circuitry **208** may be instantiated by hardware logic circuitry, which may be implemented by an ASIC or the FPGA circuitry **700** of FIG. 7 structured to perform operations corresponding to the machine readable instructions. Additionally or alternatively, the architecture-specific compiler(s) **202**, the instruction converter circuitry **204**, and/or the cost determination circuitry **208** may be instantiated by any other combination of hardware, software, and/or firmware. For example, the architecture-specific compiler(s) **202**, the instruction converter circuitry **204**, and/or the cost determination circuitry **208** may be implemented by at least one or more hardware circuits (e.g., processor circuitry, discrete and/or integrated analog and/or digital circuitry, an FPGA, an Application Specific Integrated Circuit (ASIC), a comparator, an operational-amplifier

(op-amp), a logic circuit, etc.) structured to execute some or all of the machine readable instructions and/or to perform some or all of the operations corresponding to the machine readable instructions without executing software or firmware, but other structures are likewise appropriate. (34) In some examples, the basic block execution cost determination circuitry **116** includes means for identifying, means for finding, and means for selecting, means for determining, and/or means for reporting. For example, the means for identifying may be implemented by the instruction conversion circuitry **214**, the means for finding may be implemented by the cost estimation circuitry **216**, and the means for selecting, the means for determining, and/or the means for reporting may be implemented by the calculation circuitry **218**. In some examples, the instruction conversion circuitry **214**, the cost estimation circuitry **216**, and/or the calculation circuitry **218** may be instantiated by processor circuitry such as the example processor circuitry **512** of FIG. 5. For instance, the instruction conversion circuitry **214**, the cost estimation circuitry **216**, and/or the calculation circuitry **218** may be instantiated by the example general purpose processor circuitry **600** of FIG. 6 executing machine executable instructions such as that implemented by at least blocks of FIG. 4. In some examples, the instruction conversion circuitry **214**, the cost estimation circuitry **216**, and/or the calculation circuitry **218** may be instantiated by hardware logic circuitry, which may be implemented by an ASIC or the FPGA circuitry **700** of FIG. 7 structured to perform operations corresponding to the machine readable instructions. Additionally or alternatively, the instruction conversion circuitry **214**, the cost estimation circuitry **216**, and/or the calculation circuitry **218** may be instantiated by any other combination of hardware, software, and/or firmware. For example, the instruction conversion circuitry **214**, the cost estimation circuitry **216**, and/or the calculation circuitry **218** may be implemented by at least one or more hardware circuits (e.g., processor circuitry, discrete and/or integrated analog and/or digital circuitry, an FPGA, an Application Specific Integrated Circuit (ASIC), a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to execute some or all of the machine readable instructions and/or to perform some or all of the operations corresponding to the machine readable instructions without executing software or firmware, but other structures are likewise appropriate.

(35) While an example manner of implementing the execution cost mapping table determination circuitry **104** and/or the basic block execution cost determination circuitry **116** of FIG. 1 is illustrated in FIG. 2, one or more of the elements, processes, and/or devices illustrated in FIG. 2 may be combined, divided, re-arranged, omitted, eliminated, and/or implemented in any other way. Further, the example network interface **200**, the example architecture-specific compiler(s) **202**, the example instruction converter circuitry **204**, the example cost determination circuitry **208**, the example component interface **210**, **212**, the example instruction conversion circuitry **214**, the example cost estimation circuitry **216**, the example calculation circuitry **218**, and/or, more generally, the execution cost mapping table determination circuitry **104** and/or the basic block execution cost determination circuitry **116** of FIGS. 1-2, may be implemented by hardware, software, firmware, and/or any combination of hardware, software, and/or firmware. Thus, for example, any of the example network interface **200**, the example architecture-specific compiler(s) **202**, the example instruction converter circuitry **204**, the example cost determination circuitry **208**, the example component interface **210**, **212**, the example instruction conversion circuitry **214**, the example cost estimation circuitry **216**, the example calculation circuitry **218**, and/or, more generally, the execution cost mapping table determination circuitry **104** and/or the basic block execution cost determination circuitry **116** of FIGS. 1-2, could be implemented by processor circuitry, analog circuit(s), digital circuit(s), logic circuit(s), programmable processor(s), programmable microcontroller(s), graphics processing unit(s) (GPU(s)), digital signal processor(s) (DSP(s)), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), and/or field programmable logic device(s) (FPLD(s)) such as Field Programmable Gate Arrays (FPGAs). When reading any of the apparatus or system claims of this patent to cover a purely software and/or firmware implementation, the execution cost mapping table determination

circuitry **104** and/or basic block execution cost determination circuitry **116** of FIGS. **1-2** is/are hereby expressly defined to include a non-transitory computer readable storage device or storage disk such as a memory, a digital versatile disk (DVD), a compact disk (CD), a Blu-ray disk, etc., including the software and/or firmware. Further still, the execution cost mapping table determination circuitry **104** and/or basic block execution cost determination circuitry **116** of FIGS. **1-2** may include one or more elements, processes, and/or devices in addition to, or instead of, those illustrated in FIG. **1-2**, and/or may include more than one of any or all of the illustrated elements, processes, and devices.

(36) Flowcharts representative of example hardware logic circuitry, machine readable instructions, hardware implemented state machines, and/or any combination thereof for implementing the execution cost mapping table determination circuitry **104** and/or basic block execution cost determination circuitry **116** shown in FIGS. **3A**, **3B** and/or **4**. The machine readable instructions may be one or more executable programs or portion(s) of an executable program for execution by processor circuitry, such as the processor circuitry **512** shown in the example processor platform **500** discussed below in connection with FIG. **5** and/or the example processor circuitry discussed below in connection with FIG. **5**. The program may be embodied in software stored on one or more non-transitory computer readable storage media such as a CD, a floppy disk, a hard disk drive (HDD), a DVD, a Blu-ray disk, a volatile memory (e.g., Random Access Memory (RAM) of any type, etc.), or a non-volatile memory (e.g., FLASH memory, an HDD, etc.) associated with processor circuitry located in one or more hardware devices, but the entire program and/or parts thereof could alternatively be executed by one or more hardware devices other than the processor circuitry and/or embodied in firmware or dedicated hardware. The machine readable instructions may be distributed across multiple hardware devices and/or executed by two or more hardware devices (e.g., a server and a client hardware device). For example, the client hardware device may be implemented by an endpoint client hardware device (e.g., a hardware device associated with a user) or an intermediate client hardware device (e.g., a radio access network (RAN) gateway that may facilitate communication between a server and an endpoint client hardware device). Similarly, the non-transitory computer readable storage media may include one or more mediums located in one or more hardware devices. Further, although the example program is described with reference to the flowcharts illustrated in FIGS. **3A**, **3B**, and/or **4**, many other methods of implementing the execution cost mapping table determination circuitry **104** and/or basic block execution cost determination circuitry **116** of FIGS. **1-2** may alternatively be used. For example, the order of execution of the blocks may be changed, and/or some of the blocks described may be changed, eliminated, or combined. Additionally or alternatively, any or all of the blocks may be implemented by one or more hardware circuits (e.g., processor circuitry, discrete and/or integrated analog and/or digital circuitry, an FPGA, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware. The processor circuitry may be distributed in different network locations and/or local to one or more hardware devices (e.g., a single-core processor (e.g., a single core central processor unit (CPU)), a multi-core processor (e.g., a multi-core CPU), etc.) in a single machine, multiple processors distributed across multiple servers of a server rack, multiple processors distributed across one or more server racks, a CPU and/or a FPGA located in the same package (e.g., the same integrated circuit (IC) package or in two or more separate housings, etc.).

(37) The machine readable instructions described herein may be stored in one or more of a compressed format, an encrypted format, a fragmented format, a compiled format, an executable format, a packaged format, etc. Machine readable instructions as described herein may be stored as data or a data structure (e.g., as portions of instructions, code, representations of code, etc.) that may be utilized to create, manufacture, and/or produce machine executable instructions. For example, the machine readable instructions may be fragmented and stored on one or more storage devices and/or computing devices (e.g., servers) located at the same or different locations of a

network or collection of networks (e.g., in the cloud, in edge devices, etc.). The machine readable instructions may require one or more of installation, modification, adaptation, updating, combining, supplementing, configuring, decryption, decompression, unpacking, distribution, reassignment, compilation, etc., in order to make them directly readable, interpretable, and/or executable by a computing device and/or other machine. For example, the machine readable instructions may be stored in multiple parts, which are individually compressed, encrypted, and/or stored on separate computing devices, wherein the parts when decrypted, decompressed, and/or combined form a set of machine executable instructions that implement one or more operations that may together form a program such as that described herein.

(38) In another example, the machine readable instructions may be stored in a state in which they may be read by processor circuitry, but require addition of a library (e.g., a dynamic link library (DLL)), a software development kit (SDK), an application programming interface (API), etc., in order to execute the machine readable instructions on a particular computing device or other device. In another example, the machine readable instructions may need to be configured (e.g., settings stored, data input, network addresses recorded, etc.) before the machine readable instructions and/or the corresponding program(s) can be executed in whole or in part. Thus, machine readable media, as used herein, may include machine readable instructions and/or program(s) regardless of the particular format or state of the machine readable instructions and/or program(s) when stored or otherwise at rest or in transit.

(39) The machine readable instructions described herein can be represented by any past, present, or future instruction language, scripting language, programming language, etc. For example, the machine readable instructions may be represented using any of the following languages: C, C++, Java, C#, Perl, Python, JavaScript, HyperText Markup Language (HTML), Structured Query Language (SQL), Swift, etc.

(40) As mentioned above, the example operations of FIGS. 3A, 3B and/or 4 may be implemented using executable instructions (e.g., computer and/or machine readable instructions) stored on one or more non-transitory computer and/or machine readable media such as optical storage devices, magnetic storage devices, an HDD, a flash memory, a read-only memory (ROM), a CD, a DVD, a cache, a RAM of any type, a register, and/or any other storage device or storage disk in which information is stored for any duration (e.g., for extended time periods, permanently, for brief instances, for temporarily buffering, and/or for caching of the information). As used herein, the terms non-transitory computer readable medium and non-transitory computer readable storage medium is expressly defined to include any type of computer readable storage device and/or storage disk and to exclude propagating signals and to exclude transmission media.

(41) “Including” and “comprising” (and all forms and tenses thereof) are used herein to be open ended terms. Thus, whenever a claim employs any form of “include” or “comprise” (e.g., comprises, includes, comprising, including, having, etc.) as a preamble or within a claim recitation of any kind, it is to be understood that additional elements, terms, etc., may be present without falling outside the scope of the corresponding claim or recitation. As used herein, when the phrase “at least” is used as the transition term in, for example, a preamble of a claim, it is open-ended in the same manner as the term “comprising” and “including” are open ended. The term “and/or” when used, for example, in a form such as A, B, and/or C refers to any combination or subset of A, B, C such as (1) A alone, (2) B alone, (3) C alone, (4) A with B, (5) A with C, (6) B with C, or (7) A with B and with C. As used herein in the context of describing structures, components, items, objects and/or things, the phrase “at least one of A and B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B. Similarly, as used herein in the context of describing structures, components, items, objects and/or things, the phrase “at least one of A or B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B. As used herein in the context of describing the performance or execution of processes, instructions, actions, activities

and/or steps, the phrase “at least one of A and B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B. Similarly, as used herein in the context of describing the performance or execution of processes, instructions, actions, activities and/or steps, the phrase “at least one of A or B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B.

(42) As used herein, singular references (e.g., “a”, “an”, “first”, “second”, etc.) do not exclude a plurality. The term “a” or “an” object, as used herein, refers to one or more of that object. The terms “a” (or “an”), “one or more”, and “at least one” are used interchangeably herein.

Furthermore, although individually listed, a plurality of means, elements or method actions may be implemented by, e.g., the same entity or object. Additionally, although individual features may be included in different examples or claims, these may possibly be combined, and the inclusion in different examples or claims does not imply that a combination of features is not feasible and/or advantageous.

(43) FIGS. 3A and 3B illustrate a flowchart representative of example machine readable instructions and/or example operations **300** that may be executed and/or instantiated by processor circuitry (e.g., the example execution cost mapping table determination circuitry **104** of FIG. 2) to determine develop an architecture-specific mapping table to provide to a computing device corresponding to the specific architecture. The instructions begin at block **302** when the architecture-specific compiler(s) **202** determine(s) if a program, code, and/or basic block has been obtained (e.g., via the network interface **200** for external data and/or via the component interface **210** for internal data).

(44) At block **304**, the example architecture-specific compiler(s) **202** converts the program, code, and/or basic block into assembly language. At block **306**, the example instruction converter circuitry **204** selects a first instruction (e.g., a first line of the assembly instructions) from the assembly language of the program, code, and/or basic block. At block **308**, the example instruction converter circuitry **204** converts the instruction of the assembly language into a mnemonic and an operand.

(45) At block **310**, the example instruction converter circuitry **204** determines if the operand corresponds to numeric constant(s). For example, does the operate include a numeric value or a value that corresponds to a number constant. If the example instruction converter circuitry **204** determines that the operand does not correspond to a numeric constant (block **310**: NO), control continues to block **314**. If the example instruction converter circuitry **204** determines that the operand corresponds to a numeric constant (block **310**: YES), the instruction converter circuitry **204** generates a converted instruction (e.g., an abstraction of the instruction) by replacing the number or variable corresponding to the numeric constant with a token corresponding to the constant type (e.g., integer, Boolean, short, long, float, etc.) (block **312**). In this manner, the specific numeric number is abstracted to a general number type.

(46) At block **314**, the example instruction converter circuitry **204** determines if the operand corresponds to a register operation (e.g., where a value will be accessed, stored, etc.). If the example instruction converter circuitry **204** determines that the operand does not correspond to a register (block **314**: NO), instructions continue to block **318**. If the example instruction converter circuitry **204** determines that the operand corresponds to a register (block **314**: YES), the example instruction converter circuitry **204** generates a converted instruction (e.g., an abstraction of the instruction) by replacing the value corresponding to the register with a token corresponding to register (block **316**). In this manner, the specific register number is abstracted to a general register token.

(47) At block **318**, the example instruction converter circuitry **204** pairs the instruction with the converted (e.g., abstract) instruction. At block **320**, the example instruction converter circuitry **204** determines if the pair is already stored in the example pair storage **206** (e.g., to prevent duplicate

pairs being stored). In some examples, the instruction converter circuitry **204** determines if abstraction of the instruction is already stored in the example pair storage **206** as opposed to the complete pair. If the example instruction converter circuitry **204** determines that a pair is already included in the pair storage **206** (block **320**: YES), the pair is discarded and control continues to block **322**. If the example instruction converter circuitry **204** determines that a pair is not already included in the pair storage **206** (block **320**: YES), the example instruction converter circuitry **204** stores the pair in the example pair storage **206** (block **322**).

(48) At block **324**, the example instruction converter circuitry **204** determines if the pair storage **206** includes a threshold number of pairs (e.g., entries). The less pairs stored in the example pair storage **206**, the less robust the mapping table will be (e.g., less entries and possibly less accurate execution cost information), and the more pairs, the more resources, time, and input data needed. Accordingly, a user and/or manufacturer can define the threshold of entries to balance robustness with time, resources, and/or input data. If the example pair storage **206** does not include the threshold number of pairs (e.g., does not satisfy the threshold) (block **324**: NO), control returns to block **302** to process additional programs, code and/or basic blocks. If the example pair storage **206** includes the threshold number of pairs (e.g., satisfies the threshold) (block **324**: YES), the example cost determination circuitry **208** accesses the instruction(s) corresponding to a first mnemonic (block **326** of FIG. 3B).

(49) At block **328**, the example cost determination circuitry **208** selects a first instruction of the accessed instruction(s). At block **330**, the example cost determination circuitry **208** instructs the CPU **105** to execute the selected instruction. For example, the cost determination circuitry **208** may transmit the instruction to the example CPU **105** via the component interface **210**. As described above, the example CPU **105** may execute the instruction one or more times to determine the execution cost of the instruction. After the instruction is executed one or more times, the CPU **105** provides the execution cost of the instruction to the cost determination circuitry **208** via the component interface **210**. At block **332**, the example cost mapping circuitry obtains the execution cost of the instruction from the CPU **105** via the component interface **210**.

(50) At block **334**, the example cost determination circuitry **208** determines if there is an additional instruction corresponding to the converted instruction. If the cost determination circuitry **208** determines that there is an additional instruction (block **334**: YES), the example cost determination circuitry **208** selects the subsequent instruction (block **336**) and control returns to block **330** to determine the execution cost of the additional instruction corresponding to the selected mnemonic. If the cost determination circuitry **208** determines that there is no additional instruction (block **334**: NO), the example cost determination circuitry **208** determines the execution cost average based on an average of the execution costs of the instructions corresponding to the selected mnemonic (block **338**).

(51) At block **340**, the example cost determination circuitry **208** determines the execution cost standard deviation based on the obtained execution costs of instructions corresponding to the mnemonic. At block **342**, the example cost determination circuitry **208** adds an entry for a mapping of the mnemonic to the execution cost average and standard deviation in the mapping table. For the first entry, the example cost determination circuitry **208** generates the mapping table with the single entry. For subsequent entries, the example cost mapping circuitry adds entries to the previously generated mapping table. At block **344**, the example cost determination circuitry **208** determines if there is an additional mnemonic in the pair storage **206** to be processed. If the example cost determination circuitry **208** determines that there is an additional mnemonic in the pair storage **206** (block **344**: YES), the cost determination circuitry **208** accesses the instruction(s) corresponding to the subsequent mnemonic (block **346**) and control returns to block **328**. If the example cost determination circuitry **208** determines that there is not an additional mnemonic in the pair storage **206** (block **344**: NO), the cost determination circuitry **208** deploys the mapping table to devices that correspond to the architecture via the network interface **200** (block **348**).

(52) FIG. 4 is a flowchart representative of example machine readable instructions and/or example operations that may be executed and/or instantiated by processor circuitry (e.g., the basic block execution cost determination circuitry **116** of FIG. 2) to determine execution cost of a program, code, and/or basic block using a mapping table. The instructions begin at block **402** when the example instruction conversion circuitry **214** determines if a program, code, and/or a basic block has been obtained (e.g., via the component interface **212**). It is assumed that the obtained program, code, and/or basic block is and/or has been converted to assembly language. If the obtained program, code, and/or basic block is not in assembly language, the example instruction conversion circuitry **214** converts the program, code, and/or basic block into assembly language.

(53) If the example instruction conversion circuitry **214** determines that a program, code, and/or basic block has not been obtained (block **402**: NO), control returns to block **402**. If the example instruction conversion circuitry **214** determines that a program, code, and/or basic block has been obtained (block **402**: YES), the example instruction conversion circuitry **214** selects a first instruction of the program, code, and/or basic block (block **404**). At block **406**, the example instruction conversion circuitry identifies the mnemonic and/or function of the selected instruction. At block **408**, the example cost estimation circuitry **216** accesses (e.g., via the component interface **212**) the mapping table in the mapping table storage **112** to attempt to find the mnemonic from the mapping table. If the example cost estimation circuitry **216** does not find the mnemonic and/or function in the mapping table (block **410**: NO), the example calculation circuitry **218** determines a minimum and/or maximum execution cost of the instruction based on a corresponding preset average and/or standard deviation for the instruction (block **412**). For example, the calculation circuitry **218** selects a preset average and standard deviation for the instruction and determines the minimum and/or maximum execution cost using the above Equations 1-2.

(54) If the example cost estimation circuitry **216** finds the mnemonic and/or function in the mapping table (block **410**: YES), the example calculation circuitry **218** determines a minimum and/or maximum execution cost of the instruction based on an average and/or standard deviation of the mnemonic based on the mapping table (block **414**). For example, the calculation circuitry **218** determines the average and standard deviation for the instruction based on the average and standard deviation linked to (e.g., stored in association with) the mnemonic in an entry of the mapping data and determines the minimum and/or maximum execution cost using the above Equations 1-2. At block **416**, the example instruction conversion circuitry **214** determines if there is an additional instruction in the obtained program, code, and/or basic block to be processed.

(55) If the example instruction conversion circuitry **214** determines that there is an additional instruction in the obtained program, code, and/or basic block (block **416**: YES), the example instruction conversion circuitry **214** selects the subsequent instruction (block **418**) and control returns to block **406**. If the example instruction conversion circuitry **214** determines that there is not an additional instruction in the obtained program, code, and/or basic block (block **416**: NO), the example calculation circuitry **218** determines the average and standard deviation of the executed program based on the execution cost averages and/or the minimum and maximum execution costs of the instructions of the program, code, and/or basic block (block **420**). For example, the calculation circuitry **218** determines the average and standard deviation of the execution cost for the program, code, and/or basic block using the above Equations 3-4. At block **422**, the example calculation circuitry **218** reports the execution cost to the CPU **114** and/or another device (e.g., via the component interface **212**) based on the average and standard deviation of the execution cost of the program, code, and/or basic block.

(56) FIG. 5 is a block diagram of an example processor platform **500** structured to execute and/or instantiate the machine readable instructions and/or operations of FIGS. 3A, 3B and/or 4 to implement the execution cost mapping table determination circuitry **104** and/or the basic block execution cost determination circuitry **116** of FIG. 2. The processor platform **500** can be, for example, a server, a personal computer, a workstation, a self-learning machine (e.g., a neural

network), a mobile device (e.g., a cell phone, a smart phone, a tablet such as an iPad), a personal digital assistant (PDA), an Internet appliance, a DVD player, a CD player, a digital video recorder, a Blu-ray player, a gaming console, a personal video recorder, a set top box, a headset (e.g., an augmented reality (AR) headset, a virtual reality (VR) headset, etc.) or other wearable device, or any other type of computing device.

(57) The processor platform **500** of the illustrated example includes processor circuitry **512**. The processor circuitry **512** of the illustrated example is hardware. For example, the processor circuitry **512** can be implemented by one or more integrated circuits, logic circuits, FPGAs microprocessors, CPUs, GPUs, DSPs, and/or microcontrollers from any desired family or manufacturer. The processor circuitry **512** may be implemented by one or more semiconductor based (e.g., silicon based) devices. When the processor platform **500** is implemented in the example server **102**, the processor circuitry **512** implements the architecture-specific compiler(s) **202**, the instruction converter circuitry **204**, the cost determination circuitry **208**, and/or the component interface **210**. When the processor platform **500** is implemented in the example computing device **108**, the processor circuitry **512** implements the component interface **212**, the instruction conversion circuitry **214**, the cost estimation circuitry **216**, and the calculation circuitry **218**.

(58) The processor circuitry **512** of the illustrated example includes a local memory **513** (e.g., a cache, registers, etc.). The processor circuitry **512** of the illustrated example is in communication with a main memory including a volatile memory **514** and a non-volatile memory **516** by a bus **518**. The volatile memory **514** may be implemented by Synchronous Dynamic Random Access Memory (SDRAM), Dynamic Random Access Memory (DRAM), RAMBUS® Dynamic Random Access Memory (RDRAM®), and/or any other type of RAM device. The non-volatile memory **516** may be implemented by flash memory and/or any other desired type of memory device. Access to the main memory **514**, **516** of the illustrated example is controlled by a memory controller **517**. Any one of the example memory **513**, **514**, **516** may implement the example mapping table storage **112** and/or the example pair storage **206** of FIG. 1 and/or 2.

(59) The processor platform **500** of the illustrated example also includes interface circuitry **520**. The interface circuitry **520** may be implemented by hardware in accordance with any type of interface standard, such as an Ethernet interface, a universal serial bus (USB) interface, a Bluetooth® interface, a near field communication (NFC) interface, a PCI interface, and/or a PCIe interface. In FIG. 2, the example interface **520** implements the example network interface **110** of FIG. 1 and/or the example network interface **200** of FIG. 2.

(60) In the illustrated example, one or more input devices **522** are connected to the interface circuitry **520**. The input device(s) **522** permit(s) a user to enter data and/or commands into the processor circuitry **512**. The input device(s) **522** can be implemented by, for example, an audio sensor, a microphone, a camera (still or video), a keyboard, a button, a mouse, a touchscreen, a track-pad, a trackball, an isopoint device, and/or a voice recognition system.

(61) One or more output devices **524** are also connected to the interface circuitry **520** of the illustrated example. The output devices **524** can be implemented, for example, by display devices (e.g., a light emitting diode (LED), an organic light emitting diode (OLED), a liquid crystal display (LCD), a cathode ray tube (CRT) display, an in-place switching (IPS) display, a touchscreen, etc.), a tactile output device, a printer, and/or speaker. The interface circuitry **520** of the illustrated example, thus, typically includes a graphics driver card, a graphics driver chip, and/or graphics processor circuitry such as a GPU.

(62) The interface circuitry **520** of the illustrated example also includes a communication device such as a transmitter, a receiver, a transceiver, a modem, a residential gateway, a wireless access point, and/or a network interface to facilitate exchange of data with external machines (e.g., computing devices of any kind) by a network **526**. The communication can be by, for example, an Ethernet connection, a digital subscriber line (DSL) connection, a telephone line connection, a coaxial cable system, a satellite system, a line-of-site wireless system, a cellular telephone system,

an optical connection, etc.

(63) The processor platform **500** of the illustrated example also includes one or more mass storage devices **528** to store software and/or data. Examples of such mass storage devices **528** include magnetic storage devices, optical storage devices, floppy disk drives, HDDs, CDs, Blu-ray disk drives, redundant array of independent disks (RAID) systems, solid state storage devices such as flash memory devices, and DVD drives.

(64) The machine executable instructions **532**, which may be implemented by the machine readable instructions of FIGS. **3A**, **3B** and/or **4**, may be stored in the mass storage device **528**, in the volatile memory **514**, in the non-volatile memory **516**, and/or on a removable non-transitory computer readable storage medium such as a CD or DVD.

(65) FIG. **6** is a block diagram of an example implementation of the processor circuitry **512** of FIG. **5**. In this example, the processor circuitry **512** of FIG. **5** is implemented by a microprocessor **600**. For example, the microprocessor **600** may implement multi-core hardware circuitry such as a CPU, a DSP, a GPU, an XPU, etc. Although it may include any number of example cores **602** (e.g., 1 core), the microprocessor **600** of this example is a multi-core semiconductor device including **N** cores. The cores **602** of the microprocessor **600** may operate independently or may cooperate to execute machine readable instructions. For example, machine code corresponding to a firmware program, an embedded software program, or a software program may be executed by one of the cores **602** or may be executed by multiple ones of the cores **602** at the same or different times. In some examples, the machine code corresponding to the firmware program, the embedded software program, or the software program is split into threads and executed in parallel by two or more of the cores **602**. The software program may correspond to a portion or all of the machine readable instructions and/or operations represented by the flowcharts of FIGS. **3A**, **3B**, and/or **4**.

(66) The cores **602** may communicate by an example bus **604**. In some examples, the bus **604** may implement a communication bus to effectuate communication associated with one(s) of the cores **602**. For example, the bus **604** may implement at least one of an Inter-Integrated Circuit (I2C) bus, a Serial Peripheral Interface (SPI) bus, a PCI bus, or a PCIe bus. Additionally or alternatively, the bus **604** may implement any other type of computing or electrical bus. The cores **602** may obtain data, instructions, and/or signals from one or more external devices by example interface circuitry **606**. The cores **602** may output data, instructions, and/or signals to the one or more external devices by the interface circuitry **606**. Although the cores **602** of this example include example local memory **620** (e.g., Level 1 (L1) cache that may be split into an L1 data cache and an L1 instruction cache), the microprocessor **600** also includes example shared memory **610** that may be shared by the cores (e.g., Level 2 (L2_cache)) for high-speed access to data and/or instructions. Data and/or instructions may be transferred (e.g., shared) by writing to and/or reading from the shared memory **610**. The local memory **620** of each of the cores **602** and the shared memory **610** may be part of a hierarchy of storage devices including multiple levels of cache memory and the main memory (e.g., the main memory **514**, **516** of FIG. **5**). Typically, higher levels of memory in the hierarchy exhibit lower access time and have smaller storage capacity than lower levels of memory. Changes in the various levels of the cache hierarchy are managed (e.g., coordinated) by a cache coherency policy.

(67) Each core **602** may be referred to as a CPU, DSP, GPU, etc., or any other type of hardware circuitry. Each core **602** includes control unit circuitry **614**, arithmetic, and logic (AL) circuitry (sometimes referred to as an ALU) **616**, a plurality of registers **618**, the L1 cache **620**, and an example bus **622**. Other structures may be present. For example, each core **602** may include vector unit circuitry, single instruction multiple data (SIMD) unit circuitry, load/store unit (LSU) circuitry, branch/jump unit circuitry, floating-point unit (FPU) circuitry, etc. The control unit circuitry **614** includes semiconductor-based circuits structured to control (e.g., coordinate) data movement within the corresponding core **602**. The AL circuitry **616** includes semiconductor-based circuits structured to perform one or more mathematic and/or logic operations on the data within the corresponding core **602**. The AL circuitry **616** of some examples performs integer based operations. In other

examples, the AL circuitry **616** also performs floating point operations. In yet other examples, the AL circuitry **616** may include first AL circuitry that performs integer based operations and second AL circuitry that performs floating point operations. In some examples, the AL circuitry **616** may be referred to as an Arithmetic Logic Unit (ALU). The registers **618** are semiconductor-based structures to store data and/or instructions such as results of one or more of the operations performed by the AL circuitry **616** of the corresponding core **602**. For example, the registers **618** may include vector register(s), SIMD register(s), general purpose register(s), flag register(s), segment register(s), machine specific register(s), instruction pointer register(s), control register(s), debug register(s), memory management register(s), machine check register(s), etc. The registers **618** may be arranged in a bank as shown in FIG. 6. Alternatively, the registers **618** may be organized in any other arrangement, format, or structure including distributed throughout the core **602** to shorten access time. The bus **620** may implement at least one of an I2C bus, a SPI bus, a PCI bus, or a PCIe bus

(68) Each core **602** and/or, more generally, the microprocessor **600** may include additional and/or alternate structures to those shown and described above. For example, one or more clock circuits, one or more power supplies, one or more power gates, one or more cache home agents (CHAs), one or more converged/common mesh stops (CMSs), one or more shifters (e.g., barrel shifter(s)) and/or other circuitry may be present. The microprocessor **600** is a semiconductor device fabricated to include many transistors interconnected to implement the structures described above in one or more integrated circuits (ICs) contained in one or more packages. The processor circuitry may include and/or cooperate with one or more accelerators. In some examples, accelerators are implemented by logic circuitry to perform certain tasks more quickly and/or efficiently than can be done by a general purpose processor. Examples of accelerators include ASICs and FPGAs such as those discussed herein. A GPU or other programmable device can also be an accelerator.

Accelerators may be on-board the processor circuitry, in the same chip package as the processor circuitry and/or in one or more separate packages from the processor circuitry.

(69) FIG. 7 is a block diagram of another example implementation of the processor circuitry **512** of FIG. 5. In this example, the processor circuitry **512** is implemented by FPGA circuitry **700**. The FPGA circuitry **700** can be used, for example, to perform operations that could otherwise be performed by the example microprocessor **600** of FIG. 6 executing corresponding machine readable instructions. However, once configured, the FPGA circuitry **700** instantiates the machine readable instructions in hardware and, thus, can often execute the operations faster than they could be performed by a general purpose microprocessor executing the corresponding software.

(70) More specifically, in contrast to the microprocessor **600** of FIG. 6 described above (which is a general purpose device that may be programmed to execute some or all of the machine readable instructions represented by the flowcharts of FIGS. 3A, 3B, and/or 4 but whose interconnections and logic circuitry are fixed once fabricated), the FPGA circuitry **700** of the example of FIG. 7 includes interconnections and logic circuitry that may be configured and/or interconnected in different ways after fabrication to instantiate, for example, some or all of the machine readable instructions represented by the flowcharts of FIGS. 3A, 3B, and/or 4. In particular, the FPGA **700** may be thought of as an array of logic gates, interconnections, and switches. The switches can be programmed to change how the logic gates are interconnected by the interconnections, effectively forming one or more dedicated logic circuits (unless and until the FPGA circuitry **700** is reprogrammed). The configured logic circuits enable the logic gates to cooperate in different ways to perform different operations on data received by input circuitry. Those operations may correspond to some or all of the software represented by the flowcharts of FIGS. 3A, 3B, and/or 4. As such, the FPGA circuitry **700** may be structured to effectively instantiate some or all of the machine readable instructions of the flowcharts of FIGS. 3A, 3B, and/or 4 as dedicated logic circuits to perform the operations corresponding to those software instructions in a dedicated manner analogous to an ASIC. Therefore, the FPGA circuitry **700** may perform the operations

corresponding to the some or all of the machine readable instructions of FIG. 7 faster than the general purpose microprocessor can execute the same.

(71) In the example of FIG. 7, the FPGA circuitry **700** is structured to be programmed (and/or reprogrammed one or more times) by an end user by a hardware description language (HDL) such as Verilog. The FPGA circuitry **700** of FIG. 7, includes example input/output (I/O) circuitry **702** to obtain and/or output data to/from example configuration circuitry **704** and/or external hardware (e.g., external hardware circuitry) **706**. For example, the configuration circuitry **704** may implement interface circuitry that may obtain machine readable instructions to configure the FPGA circuitry **700**, or portion(s) thereof. In some such examples, the configuration circuitry **704** may obtain the machine readable instructions from a user, a machine (e.g., hardware circuitry (e.g., programmed, or dedicated circuitry) that may implement an Artificial Intelligence/Machine Learning (AI/ML) model to generate the instructions), etc. In some examples, the external hardware **706** may implement the microprocessor **600** of FIG. 6. The FPGA circuitry **700** also includes an array of example logic gate circuitry **708**, a plurality of example configurable interconnections **710**, and example storage circuitry **712**. The logic gate circuitry **708** and interconnections **710** are configurable to instantiate one or more operations that may correspond to at least some of the machine readable instructions of FIG. 4-5 and/or other desired operations. The logic gate circuitry **708** shown in FIG. 7 is fabricated in groups or blocks. Each block includes semiconductor-based electrical structures that may be configured into logic circuits. In some examples, the electrical structures include logic gates (e.g., And gates, Or gates, Nor gates, etc.) that provide basic building blocks for logic circuits. Electrically controllable switches (e.g., transistors) are present within each of the logic gate circuitry **708** to enable configuration of the electrical structures and/or the logic gates to form circuits to perform desired operations. The logic gate circuitry **708** may include other electrical structures such as look-up tables (LUTs), registers (e.g., flip-flops or latches), multiplexers, etc.

(72) The interconnections **710** of the illustrated example are conductive pathways, traces, vias, or the like that may include electrically controllable switches (e.g., transistors) whose state can be changed by programming (e.g., using an HDL instruction language) to activate or deactivate one or more connections between one or more of the logic gate circuitry **708** to program desired logic circuits.

(73) The storage circuitry **712** of the illustrated example is structured to store result(s) of the one or more of the operations performed by corresponding logic gates. The storage circuitry **712** may be implemented by registers or the like. In the illustrated example, the storage circuitry **712** is distributed amongst the logic gate circuitry **708** to facilitate access and increase execution speed.

(74) The example FPGA circuitry **700** of FIG. 7 also includes example Dedicated Operations Circuitry **714**. In this example, the Dedicated Operations Circuitry **714** includes special purpose circuitry **716** that may be invoked to implement commonly used functions to avoid the need to program those functions in the field. Examples of such special purpose circuitry **716** include memory (e.g., DRAM) controller circuitry, PCIe controller circuitry, clock circuitry, transceiver circuitry, memory, and multiplier-accumulator circuitry. Other types of special purpose circuitry may be present. In some examples, the FPGA circuitry **700** may also include example general purpose programmable circuitry **718** such as an example CPU **720** and/or an example DSP **722**. Other general purpose programmable circuitry **718** may additionally or alternatively be present such as a GPU, an XPU, etc., that can be programmed to perform other operations.

(75) Although FIGS. 6 and 7 illustrate two example implementations of the processor circuitry **512** of FIG. 5, many other approaches are contemplated. For example, as mentioned above, modern FPGA circuitry may include an on-board CPU, such as one or more of the example CPU **720** of FIG. 7. Therefore, the processor circuitry **512** of FIG. 5 may additionally be implemented by combining the example microprocessor **600** of FIG. 6 and the example FPGA circuitry **700** of FIG. 7. In some such hybrid examples, a first portion of the machine readable instructions represented

by the flowcharts of FIGS. 3A, 3B, and/or 4 may be executed by one or more of the cores 602 of FIG. 6 and a second portion of the machine readable instructions represented by the flowcharts of FIGS. 3A, 3B, and/or 4 may be executed by the FPGA circuitry 700 of FIG. 7.

(76) In some examples, the processor circuitry 512 of FIG. 5 may be in one or more packages. For example, the processor circuitry 600 of FIG. 6 and/or the FPGA circuitry 700 of FIG. 7 may be in one or more packages. In some examples, an XPU may be implemented by the processor circuitry 512 of FIG. 5, which may be in one or more packages. For example, the XPU may include a CPU in one package, a DSP in another package, a GPU in yet another package, and an FPGA in still yet another package.

(77) A block diagram illustrating an example software distribution platform 805 to distribute software such as the example machine readable instructions 532 of FIG. 5 to hardware devices owned and/or operated by third parties is illustrated in FIG. 8. The example software distribution platform 805 may be implemented by any computer server, data facility, cloud service, etc., capable of storing and transmitting software to other computing devices. The third parties may be customers of the entity owning and/or operating the software distribution platform 805. For example, the entity that owns and/or operates the software distribution platform 805 may be a developer, a seller, and/or a licensor of software such as the example machine readable instructions 532 of FIG. 5. The third parties may be consumers, users, retailers, OEMs, etc., who purchase and/or license the software for use and/or re-sale and/or sub-licensing. In the illustrated example, the software distribution platform 805 includes one or more servers and one or more storage devices. The storage devices store the machine readable instructions 532, which may correspond to the example machine readable instructions 300, 400 of FIGS. 3A, 3B and/or 4, as described above. The one or more servers of the example software distribution platform 805 are in communication with a network 810, which may correspond to any one or more of the Internet and/or any example network. In some examples, the one or more servers are responsive to requests to transmit the software to a requesting party as part of a commercial transaction. Payment for the delivery, sale, and/or license of the software may be handled by the one or more servers of the software distribution platform and/or by a third party payment entity. The servers enable purchasers and/or licensors to download the machine readable instructions 532 from the software distribution platform 805. For example, the software, which may correspond to the example machine readable instructions 300, 400 of FIGS. 3A, 3B and/or 4, may be downloaded to the example processor platform 500, which is to execute the machine readable instructions 532 to implement the execution cost mapping table determination circuitry 104. In some example, one or more servers of the software distribution platform 805 periodically offer, transmit, and/or force updates to the software (e.g., the example machine readable instructions 532 of FIG. 5) to ensure improvements, patches, updates, etc., are distributed and applied to the software at the end user devices.

(78) Example methods, apparatus, systems, and articles of manufacture to improve data quality for artificial intelligence are disclosed herein. Further examples and combinations thereof include the following: Example 1 includes an apparatus to generate a mapping table based on execution cost, the apparatus comprising memory, instructions included in the apparatus, and processor circuitry to execute the instructions to cause a plurality of instructions corresponding to a mnemonic to be executed, determine an average execution cost of the plurality of instructions, determine a standard deviation of execution costs of the plurality of instructions, and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

(79) Example 2 includes the apparatus of example 1, wherein the processor circuitry is to cause the plurality of instructions to be executed on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

(80) Example 3 includes the apparatus of example 1, wherein the plurality of instructions is a first

plurality of instructions, the processor circuitry is to convert a second plurality of instructions into mnemonics and operands, generate a converted instructions by replacing the operands with a token corresponding to the operands, and generate pairs by combining the instructions with the converted instructions.

(81) Example 4 includes the apparatus of example 3, wherein the processor circuitry is to store the pairs in the memory, and access the first plurality of instruction corresponding to the mnemonic from the memory.

(82) Example 5 includes the apparatus of example 3, wherein the processor circuitry is to determine whether a pair including at least one of a same mnemonic or a same operand is included in storage, and when the pair is already included in the storage, discard the pair.

(83) Example 6 includes the apparatus of example 3, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

(84) Example 7 includes the apparatus of example 1, wherein the processor circuitry is to obtain the instructions from a repository, and convert the instructions into assembly language.

(85) Example 8 includes a non-transitory computer readable medium comprising instructions which, when executed, cause one or more processors to at least cause a plurality of instructions corresponding to a mnemonic to be executed, determine an average execution cost of the plurality of instructions, determine a standard deviation of execution costs of the plurality of instructions, and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

(86) Example 9 includes the computer readable storage medium of example 8, wherein the instructions cause the one or more processors to cause the plurality of instructions to be executed on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

(87) Example 10 includes the computer readable storage medium of example 8, wherein the plurality of instructions is a first plurality of instructions, the instructions to cause the one or more processors to convert a second plurality of instructions into mnemonics and operands, generate a converted instructions by replacing the operands with a token corresponding to the operands, and generate pairs by combining the instructions with the converted instructions.

(88) Example 11 includes the computer readable storage medium of example 10, wherein the instructions cause the one or more processors to store the pairs in storage, and access the first plurality of instruction corresponding to the mnemonic from the storage.

(89) Example 12 includes the computer readable storage medium of example 10, wherein the instructions cause the one or more processors to determine whether a pair including at least one of a same mnemonic or a same operand is included in storage, and when the pair is already included in the storage, discard the pair.

(90) Example 13 includes the computer readable storage medium of example 10, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

(91) Example 14 includes the computer readable storage medium of example 8, wherein the instructions cause the one or more processors to obtain the instructions from a repository, and convert the instructions into assembly language.

(92) Example 15 includes an apparatus to generate a mapping table based on execution cost, the apparatus comprising interface circuitry, and processor circuitry including one or more of at least one of a central processing unit, a graphic processing unit or a digital signal processor, the at least one of the central processing unit, the graphic processing unit or the digital signal processor having control circuitry, one or more registers, and arithmetic and logic circuitry to perform one or more first operations corresponding to instructions in the apparatus, and, a Field Programmable Gate Array (FPGA), the FPGA including logic gate circuitry, a plurality of configurable

interconnections, and storage circuitry, the logic gate circuitry and interconnections to perform one or more second operations, or Application Specific Integrate Circuitry (ASIC) including logic gate circuitry to perform one or more third operations, the processor circuitry to perform at least one of the first operations, the second operations or the third operations to instantiate cost mapping circuitry to cause a plurality of instructions corresponding to a mnemonic to be executed, determine an average execution cost of the plurality of instructions, determine a standard deviation of execution costs of the plurality of instructions, and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

(93) Example 16 includes the apparatus of example 15, wherein the cost mapping circuitry is to cause the plurality of instructions to be executed on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

(94) Example 17 includes the apparatus of example 15, wherein the plurality of instructions is a first plurality of instructions, further including instruction converter circuitry is to convert a second plurality of instructions into mnemonics and operands, generate a converted instructions by replacing the operands with a token corresponding to the operands, and generate pairs by combining the instructions with the converted instructions.

(95) Example 18 includes the apparatus of example 17, wherein the instruction converter circuitry is to store the pairs in the storage, and the cost mapping circuitry is to access the first plurality of instruction corresponding to the mnemonic from the storage.

(96) Example 19 includes the apparatus of example 17, wherein the cost mapping circuitry is to determine whether a pair including at least one of a same mnemonic or a same operand is included in storage, and when the pair is already included in the storage, discard the pair.

(97) Example 20 includes the apparatus of example 17, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

(98) Example 21 includes the apparatus of example 15, further including an architecture-specific compiler to obtain the instructions from a repository, and convert the instructions into assembly language.

(99) Example 22 includes an apparatus to generate a mapping table based on execution cost, the apparatus comprising means for identify a mnemonic of a plurality of instructions, means for generating a mapping table, the means for generating to cause the plurality of instructions corresponding to the mnemonic to be executed, determine an average execution cost of the plurality of instructions, determine standard deviation of execution costs of the plurality of instructions, and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

(100) Example 23 includes the apparatus of example 22, wherein the means for generating is to cause the plurality of instructions to be executed on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

(101) Example 24 includes the apparatus of example 22, wherein the plurality of instructions is a first plurality of instructions, further including means for converting, the means for converting to convert a second plurality of instructions into mnemonics and operands, generate a converted instructions by replacing the operands with a token corresponding to the operands, and generate pairs by combining the instructions with the converted instructions.

(102) Example 25 includes the apparatus of example 24, wherein the means for converting is to store the pairs in storage, and access the first plurality of instruction corresponding to the mnemonic from the storage.

(103) Example 26 includes the apparatus of example 24, wherein the means for generating is to determine whether a pair including at least one of a same mnemonic or a same operand is included

in storage, and when the pair is already included in the storage, discard the pair.

(104) Example 27 includes the apparatus of example 24, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

(105) Example 28 includes the apparatus of example 22, further including means for converting, the means for converting to obtain the instructions from a repository, and convert the instructions into assembly language.

(106) Example 29 includes a method to generate a mapping table based on execution cost, the method comprising causing a plurality of instructions corresponding to a mnemonic to be executed, determining, by executing an instruction with one or more processors, an average execution cost of the plurality of instructions, determining, by executing an instruction with the one or more processors, a standard deviation of execution costs of the plurality of instructions, and generating, by executing an instruction with the one or more processors, a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

(107) Example 30 includes the method of example 29, further including causing the plurality of instructions to be executed on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

(108) Example 31 includes the method of example 29, wherein the plurality of instructions is a first plurality of instructions, further including converting a second plurality of instructions into mnemonics and operands, generating a converted instructions by replacing the operands with a token corresponding to the operands, and generating pairs by combining the instructions with the converted instructions.

(109) Example 32 includes the method of example 31, further including storing the pairs in the storage, and accessing the first plurality of instruction corresponding to the mnemonic from the storage.

(110) Example 33 includes the method of example 31, further including determining whether a pair including at least one of a same mnemonic or a same operand is included in storage, and when the pair is already included in the storage, discarding the pair.

(111) Example 34 includes the method of example 31, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

(112) Example 35 includes the method of example 29, further including obtaining the instructions from a repository, and converting the instructions into assembly language.

(113) From the foregoing, it will be appreciated that example systems, methods, apparatus, and articles of manufacture have been disclosed that determine execution cost. To determine the execution cost of instructions, some techniques run the instructions prior to executing the instructions to identify the code execution. However, such techniques are expensive and time consuming. Other techniques have been used to estimate execution cost using models and/or other tools. However, such techniques may be inaccurate and/or require large maintenance overhead. Some techniques use artificial intelligence (AI) to predict the execution cost of a set of instructions. Although such techniques are accurate, such techniques are computationally intensive and AI-based techniques do not provide an explanation for a result.

(114) Examples disclosed herein include a server that develops architecture-specific mapping tables that maps converted instructions (e.g., abstractions of instructions based on the leading mnemonic and/or function of an instruction line) to an average execution cost and standard deviation cost. Using examples disclosed herein, the server can obtain programs and/or basic blocks of code from storage (e.g., open-source repositories, crowd-sourced repositories, closed-source repositories, local storage, external stored, local databases, external databases, etc.), execute lines of the code to identify execution costs corresponding to the function and/or mnemonic of the line of code and map the mnemonic to the average execution cost and the standard deviation of the execution cost.

Examples disclosed herein develops an approach that is accurate, explainable, and executes using less processor resources than the above-mentioned techniques. Thus, the disclosed systems, methods, apparatus, and articles of manufacture are accordingly directed to one or more improvement(s) in the operation of a machine such as a computer or other electronic and/or mechanical device.

(115) Although certain example systems, methods, apparatus, and articles of manufacture have been disclosed herein, the scope of coverage of this patent is not limited thereto. On the contrary, this patent covers all systems, methods, apparatus, and articles of manufacture fairly falling within the scope of the claims of this patent.

(116) The following claims are hereby incorporated into this Detailed Description by this reference, with each claim standing on its own as a separate embodiment of the present disclosure.

Claims

1. An apparatus to generate a mapping table based on execution cost, the apparatus comprising: memory; instructions included in the apparatus; and processor circuitry to execute the instructions to: convert a plurality of instructions into mnemonics and operands; generate converted instructions by replacing the operands with corresponding tokens; generate unique pairs by combining respective ones of the plurality of instructions with the corresponding converted instructions; execute a portion of the plurality of instructions corresponding to a mnemonic, the portion of the plurality of instructions selected from the generated pairs; determine an average execution cost of the portion of the plurality of instructions; determine a standard deviation of execution costs of the portion of the plurality of instructions; and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.
2. The apparatus of claim 1, wherein the processor circuitry is to execute the portion of the plurality of instructions on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.
3. The apparatus of claim 1, wherein the processor circuitry is to: store the unique pairs in the memory; and access the portion of the plurality of instructions corresponding to the mnemonic from the memory.
4. The apparatus of claim 1, wherein the processor circuitry is to: determine whether a pair including at least one of a same mnemonic or a same operand is included in storage; and when the pair is already included in the storage, discard the pair.
5. The apparatus of claim 1, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.
6. The apparatus of claim 1, wherein the processor circuitry is to: obtain the instructions from a repository; and convert the instructions into assembly language.
7. A non-transitory computer readable medium comprising instructions which, when executed, cause one or more processors to at least: convert a plurality of instructions into mnemonics and operands; generate converted instructions by replacing the operands with corresponding tokens; generate unique pairs by combining respective ones of the plurality of instructions with the corresponding converted instructions; execute a portion of the plurality of instructions corresponding to a mnemonic, the portion of the plurality of instructions selected from the generated pairs; determine an average execution cost of the portion of the plurality of instructions; determine a standard deviation of execution costs of the portion of the plurality of instructions; and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.
8. The computer readable medium of claim 7, wherein the instructions cause the one or more processors to execute the portion of the plurality of instructions on a processor unit specific to an

architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

9. The computer readable medium of claim 8, wherein the instructions cause the one or more processors to: store the unique pairs in storage; and access the portion of the plurality of instructions corresponding to the mnemonic from the storage.

10. The computer readable storage medium of claim 8, wherein the instructions cause the one or more processors to: determine whether a pair including at least one of a same mnemonic or a same operand is included in storage; and when the pair is already included in the storage, discard the pair.

11. The computer readable medium of claim 8, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

12. The computer readable medium of claim 7, wherein the instructions cause the one or more processors to: obtain the instructions from a repository; and convert the instructions into assembly language.

13. An apparatus to generate a mapping table based on execution cost, the apparatus comprising: interface circuitry; and processor circuitry including one or more of: at least one of a central processing unit, a graphic processing unit or a digital signal processor, the at least one of the central processing unit, the graphic processing unit or the digital signal processor having control circuitry, one or more registers, and arithmetic and logic circuitry to perform one or more first operations corresponding to instructions in the apparatus, and; a Field Programmable Gate Array (FPGA), the FPGA including logic gate circuitry, a plurality of configurable interconnections, and storage circuitry, the logic gate circuitry and interconnections to perform one or more second operations; or Application Specific Integrate Circuitry (ASIC) including logic gate circuitry to perform one or more third operations; the processor circuitry to perform at least one of the first operations, the second operations or the third operations to instantiate: instruction converter circuitry to: convert a plurality of instructions into mnemonics and operands; generate converted instructions by replacing the operands with corresponding tokens; generate unique pairs by combining respective ones of the plurality of instructions with the corresponding converted instructions; cost mapping circuitry to: execute a portion of the plurality of instructions corresponding to a mnemonic, the portion of the plurality of instructions selected from the generated pairs; determine an average execution cost of the portion of the plurality of instructions; determine a standard deviation of execution costs of the portion of the plurality of instructions; and generate a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

14. The apparatus of claim 13, wherein the cost mapping circuitry is to execute the portion of the plurality of instructions on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

15. The apparatus of claim 13, wherein: the instruction converter circuitry is to store the unique pairs in storage; and the cost mapping circuitry is to access the portion of the plurality of instructions corresponding to the mnemonic from the storage.

16. The apparatus of claim 13, wherein the cost mapping circuitry is to: determine whether a pair including at least one of a same mnemonic or a same operand is included in storage; and when the pair is already included in the storage, discard the pair.

17. The apparatus of claim 13, wherein the token corresponds to at least one of a constant type when the operand corresponds to a numeric constant or a register when the operand corresponds to a register operation.

18. The apparatus of claim 13, further including an architecture-specific compiler to: obtain the instructions from a repository; and convert the instructions into assembly language.

19. A method to generate a mapping table based on execution cost, the method comprising: converting, by executing an instruction with one or more processors, a plurality of instructions into

mnemonics and operands; generating, by executing an instruction with the one or more processors, converted instructions by replacing the operands with corresponding tokens; generating, by executing an instruction with the one or more processors, unique pairs by combining respective ones of the plurality of instructions with the corresponding converted instructions; executing a portion of the plurality of instructions corresponding to a mnemonic, the portion of the plurality of instructions selected from the generated pairs; determining, by executing an instruction with the one or more processors, an average execution cost of the portion of the plurality of instructions; determining, by executing an instruction with the one or more processors, a standard deviation of execution costs of the portion of the plurality of instructions; and generating, by executing an instruction with the one or more processors, a mapping table including an entry, the entry including the mnemonic in association with the average and the standard deviation.

20. The method of claim 19, further including executing the portion of the plurality of instructions on a processor unit specific to an architecture, the architecture corresponding to a computing device that obtains the mapping table to estimate execution cost.

21. The method of claim 19, further including: storing the unique pairs in storage; and accessing the portion of the plurality of instructions corresponding to the mnemonic from the storage.
