

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386678
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Brener-Shalem; Irit et al.

Intelligent load scheduling in a storage system

Abstract

Techniques are provided for implementing intelligent load scheduling in a storage system. In an illustrative embodiment, a method comprises collecting a plurality of metrics corresponding to one or more workloads, and processing the plurality of metrics using a machine learning model to predict a load on at least one server. The machine learning model predicts the load on the at least one server based, at least in part, on weights for respective ones of the plurality of metrics generated during training of the machine learning model. In the method, execution by the at least one server of one or more processes in addition to the one or more workloads is scheduled based, at least in part, on the load predicted by the machine learning model.

Inventors: Brener-Shalem; Irit (Ramat Gan, IL), Yarimi; Michal (Rehovot, IL), Yeshanov; Yonatan (Tel Aviv, IL)

Applicant: Dell Products L.P. (Round Rock, TX)

Family ID: 88415430

Assignee: Dell Products L.P. (Round Rock, TX)

Appl. No.: 17/726853

Filed: April 22, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230342221 A1	Oct. 26, 2023

Publication Classification

Int. Cl.: G06F9/50 (20060101); G06N20/00 (20190101)

U.S. Cl.:

Field of Classification Search

CPC: G06F (9/5083)

References Cited**U.S. PATENT DOCUMENTS**

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5381539	12/1994	Yanai et al.	N/A	N/A
5551003	12/1995	Mattson et al.	N/A	N/A
5764880	12/1997	Gerdt et al.	N/A	N/A
6052799	12/1999	Li et al.	N/A	N/A
6941420	12/2004	Butterworth et al.	N/A	N/A
8843676	12/2013	Rajamanickam et al.	N/A	N/A
9372751	12/2015	McNutt	N/A	N/A
9514014	12/2015	Webman et al.	N/A	N/A
9892045	12/2017	Douglis et al.	N/A	N/A
10078598	12/2017	Wallace et al.	N/A	N/A
10331561	12/2018	Shilane et al.	N/A	N/A
10445180	12/2018	Butterworth et al.	N/A	N/A
10986174	12/2020	Sharma et al.	N/A	N/A
11119668	12/2020	Keller et al.	N/A	N/A
11144399	12/2020	Yarimi et al.	N/A	N/A
11163479	12/2020	Liebllich et al.	N/A	N/A
11163699	12/2020	Keller et al.	N/A	N/A
11221975	12/2021	Puder et al.	N/A	N/A
11262933	12/2021	Matosevich et al.	N/A	N/A
11301162	12/2021	Matosevich et al.	N/A	N/A
11307935	12/2021	Keller et al.	N/A	N/A
11740933	12/2022	Vega	718/102	G06F 16/9024
2002/0032835	12/2001	Li et al.	N/A	N/A
2008/0021853	12/2007	Modha et al.	N/A	N/A
2009/0204761	12/2008	Caprioli et al.	N/A	N/A
2009/0276593	12/2008	Jacobson et al.	N/A	N/A
2011/0276538	12/2010	Knapp	709/219	G06F 16/1827
2013/0305002	12/2012	Hallak et al.	N/A	N/A
2014/0053151	12/2013	Heninger	718/1	G06F 9/5077
2014/0165071	12/2013	Celis	718/104	G06F 9/5066
2014/0215147	12/2013	Pan	N/A	N/A
2014/0215262	12/2013	Li et al.	N/A	N/A
2014/0244935	12/2013	Ezra et al.	N/A	N/A
2016/0103764	12/2015	Banerjee et al.	N/A	N/A
2018/0113640	12/2017	Fernandez et al.	N/A	N/A
2018/0267893	12/2017	Barzik et al.	N/A	N/A
2018/0300075	12/2017	Fernandez et al.	N/A	N/A

2019/0163587	12/2018	Anna et al.	N/A	N/A
2019/0227845	12/2018	Sridhar et al.	N/A	N/A
2020/0133503	12/2019	Sun et al.	N/A	N/A
2021/0279187	12/2020	Puder et al.	N/A	N/A
2021/0294505	12/2020	Keller et al.	N/A	N/A
2021/0294774	12/2020	Keller et al.	N/A	N/A
2021/0294775	12/2020	Keller et al.	N/A	N/A
2021/0303160	12/2020	Lieblich et al.	N/A	N/A
2021/0303169	12/2020	Tagar et al.	N/A	N/A
2021/0303202	12/2020	Ben Zeev et al.	N/A	N/A
2021/0303401	12/2020	Yarimi et al.	N/A	N/A
2021/0303407	12/2020	Keller et al.	N/A	N/A
2021/0303480	12/2020	Keller et al.	N/A	N/A
2021/0373796	12/2020	Matosevich et al.	N/A	N/A
2022/0004320	12/2021	Matosevich et al.	N/A	N/A
2022/0035788	12/2021	Aharoni et al.	N/A	N/A
2022/0113867	12/2021	Aharoni et al.	N/A	N/A
2022/0114184	12/2021	Sharma et al.	N/A	N/A
2022/0116454	12/2021	Aharoni et al.	N/A	N/A
2022/0121458	12/2021	Moran et al.	N/A	N/A
2023/0089593	12/2022	O'Reilly	718/104	G06F 1/3228
2023/0281472	12/2022	Asghari	706/12	G06N 20/20

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2015108670	12/2019	WO	N/A
2020204880	12/2019	WO	N/A
2020204882	12/2019	WO	N/A

OTHER PUBLICATIONS

Bodas, Deva, et al. "Simple power-aware scheduler to limit power consumption by hpc system within a budget." 2014 Energy Efficient Supercomputing Workshop. IEEE, 2014. (Year: 2014). cited by examiner

International Search Report and Written Opinion of PCT/US2019/024885 dated Jan. 7, 2020, 13 pages. cited by applicant

International Search Report and Written Opinion of PCT/US2019/024900 dated Jan. 7, 2020, 12 pages. cited by applicant

Dell EMC, "EMC ScaleIO Basic Architecture Documentation," Technical White Paper, Mar. 2017, 22 pages. cited by applicant

EMC2, "EMC ScaleIO Design Considerations and Best Practices," Technical White Paper, Jun. 2016, 30 pages. cited by applicant

I. Koltsidas et al., "SoftwAre Log-Structured Array (SALSA)—A Unified Stack for SSDs and SMR Disks," IBM Research Report, Dec. 2, 2015, 13 pages. cited by applicant

S. M. Rumble et al., "Log-Structured Memory for DRAM-Based Storage," Proceedings of the 12th USENIX Conference on File and Storage Technologies, Santa Clara, CA, Feb. 17-20, 2014, 17 pages. cited by applicant

Dell EMC, "Dell EMC VxFlex Family Overview," Technical White Paper, May 2019, 44 pages. cited by applicant

J. Nakano et al., "ReViveI/O: Efficient Handling of I/O in Highly-Available Rollback-Recovery Servers," HPCA, 10.1109/2006.1598129, pp. 200-211. cited by applicant

Wikipedia, "Raft (Computer Science)," [https://en.wikipedia.org/wiki/Raft_\(computer_science\)](https://en.wikipedia.org/wiki/Raft_(computer_science)), Feb. 10, 2020, 4 pages. cited by applicant

Wikipedia, "Paxos (Computer Science)," [https://en.wikipedia.org/wiki/Paxos_\(computer_science\)](https://en.wikipedia.org/wiki/Paxos_(computer_science)), Dec. 6, 2019, 21 pages. cited by applicant

Wikipedia, "State Machine Replication," https://en.wikipedia.org/wiki/State_machine_replication, Dec. 14, 2019, 9 pages. cited by applicant

Dell Technologies, "Dell EMC PowerFlex: Secure Snapshots," Technical White Paper, Jul. 2020, 17 pages. cited by applicant

Dell Technologies, "Dell EMC PowerFlex: Protected Maintenance Mode," Technical White Paper, Jul. 2020, 20 pages. cited by applicant

Dell Technologies, "Dell EMC PowerFlex: Introduction to Replication," Technical White Paper, Jun. 2020, 34 pages. cited by applicant

Dell Technologies, "Dell EMC PowerFlex: Networking Best Practices and Design Considerations," Best Practices, Jun. 2020, 64 pages. cited by applicant

Dell EMC, "Getting To Know Dell EMC PowerFlex," Version 3.5.x, Rev. 02, Jan. 2021, 66 pages. cited by applicant

Dell EMC, "Dell EMC VxRack FLEX," Dell EMC Product Overview, 2018, 5 pages. cited by applicant

G. Soundararajan et al., "Dynamic Resource Allocation for Database Servers Running on Virtual Storage," FAST 2009: Proceedings of the 7th conference on File and storage technologies, Feb. 2009, pp. 71-84. cited by applicant

T. McAleer, "Interpreting Linear Regression Through statsmodels .summary()" <https://medium.com/swlh/interpreting-linear-regression-through-statsmodels-summary-4796d359035a>, Dec. 5, 2020, 7 pages. cited by applicant

U.S. Appl. No. 17/077,105 filed in the name of Yosef Shatsky et al. Oct. 22, 2020, and entitled "Volume Tiering in Storage Systems." cited by applicant

U.S. Appl. No. 17/106,988 filed in the name of Rivka Matosevich et al. Nov. 30, 2020, and entitled "Managing Host Connectivity to a Data Storage System." cited by applicant

U.S. Appl. No. 17/123,525 filed in the name of Itay Keller et al. Dec. 16, 2020, and entitled "Deferred Reclamation of Invalidated Entries Associated with Replication in a Log-Structured Array." cited by applicant

U.S. Appl. No. 17/145,646 filed in the name of Yosef Shatsky et al. Jan. 11, 2021, and entitled "Redistribution of Processing Groups between Server Nodes Based on Hardware Resource Utilization." cited by applicant

U.S. Appl. No. 17/232,203 filed in the name of Roman Spiegelman Apr. 16, 2021, and entitled "Object Synchronization of Server Nodes in a Network Computing Environment." cited by applicant

U.S. Appl. No. 17/236,256 filed in the name of Doron Tal et al. Apr. 21, 2021, and entitled "Recovery from Partial Device Error in Data Storage System." cited by applicant

U.S. Appl. No. 17/306,601 filed in the name of Rivka Matosevich et al. May 3, 2021, and entitled "Managing Replication Journal in a Distributed Replication System." cited by applicant

U.S. Appl. No. 17/308,166 filed in the name of Adi Bar Shalom et al. May 5, 2021, and entitled "Journal Barrier Consistency Determination." cited by applicant

U.S. Appl. No. 17/351,733 filed in the name of Yosef Shatsky et al. Jun. 18, 2021, and entitled "Data Deduplication in a Disaggregated Storage System." cited by applicant

U.S. Appl. No. 17/361,666 filed in the name of Yosef Shatsky et al. Jun. 29, 2021, and entitled "Tracking Utilization of Data Blocks in a Storage System." cited by applicant

U.S. Appl. No. 17/511,695 filed in the name of Yosef Shatsky et al. Oct. 27, 2021, and entitled "Write Cache Management." cited by applicant

U.S. Appl. No. 17/511,699 filed in the name of Yosef Shatsky et al. Oct. 27, 2021, and entitled

“Metadata Management in Storage Systems.” cited by applicant
U.S. Appl. No. 17/512,890 filed in the name of Yosef Shatsky et al. Oct. 28, 2021, and entitled
“Utilizing Checkpoints for Resiliency of Metadata in Storage Systems.” cited by applicant
U.S. Appl. No. 17/569,198 filed in the name of Yosef Shatsky Jan. 5, 2022, and entitled “Utilizing
a Persistent Write Cache as a Redo Log.” cited by applicant
U.S. Appl. No. 17/583,365 filed in the name of Doron Tal et al. Jan. 25, 2022, and entitled “Data
Deduplication in a Storage System.” cited by applicant
U.S. Appl. No. 17/583,787 filed in the name of Michal Yarimi et al. Jan. 25, 2022, and entitled
“Intelligent Defragmentation in a Storage System.” cited by applicant

Primary Examiner: Sun; Charlie

Attorney, Agent or Firm: Ryan, Mason & Lewis, LLP

Background/Summary

FIELD

(1) The field relates generally to information processing systems, and more particularly to storage in information processing systems.

BACKGROUND

(2) A wide variety of different types of distributed storage systems are known. Such storage systems include clustered storage systems as well as other types of storage systems that are distributed across multiple storage nodes. Distributed storage systems can include a potentially large number of distributed storage nodes that are interconnected by a mesh network or other type of communication network. Each such storage node of a distributed storage system typically processes input-output (IO) operations from one or more host devices and in processing those TO operations runs various storage application processes that generally involve interaction of that storage node with one or more other ones of the storage nodes.

(3) In distributed storage systems, workloads (e.g., TO operations) from host devices may be distributed among multiple servers (e.g., storage nodes) to provide, for example, load balancing, high availability, utilization of resources available from different servers and/or fault tolerance. To ensure adequate performance levels, the percentage of a workload performed by a given server should be proportional to the server's processing resources and capabilities.

SUMMARY

(4) Exemplary embodiments of the disclosure include techniques for implementing intelligent load scheduling in a storage system. For example, an exemplary embodiment includes a method which comprises collecting a plurality of metrics corresponding to one or more workloads, and processing the plurality of metrics using a machine learning model to predict a load on at least one server. The machine learning model predicts the load on the at least one server based, at least in part, on weights for respective ones of the plurality of metrics generated during training of the machine learning model. In the method, execution by the at least one server of one or more processes in addition to the one or more workloads is scheduled based, at least in part, on the load predicted by the machine learning model.

(5) Other embodiments of the disclosure include, without limitation, systems and articles of manufacture comprising processor-readable storage media, which are configured to implement intelligent load scheduling in a storage system.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 schematically illustrates a network computing system comprising a storage system incorporating functionality for intelligent load scheduling in an illustrative embodiment.
- (2) FIG. 2 schematically illustrates a storage node which implements an intelligent load scheduling system in an illustrative embodiment.
- (3) FIG. 3 depicts a plot of an output of multiple linear regression on metrics processed using principal component analysis (PCA) in an illustrative embodiment.
- (4) FIG. 4 depicts a table of an output of multiple linear regression on metrics processed using PCA in an illustrative embodiment.
- (5) FIG. 5 depicts a plot of an output of multiple linear regression on metrics that are different from those used in connection with FIGS. 3 and 4 and processed using PCA in an illustrative embodiment.
- (6) FIG. 6 depicts an alternative plot of the output of multiple linear regression on the metrics of FIG. 5 in an illustrative embodiment.
- (7) FIG. 7 illustrates a flow diagram of a method for intelligent load scheduling in an illustrative embodiment.
- (8) FIG. 8 schematically illustrates a framework of a server for hosting a storage control system which implements an intelligent load scheduling system in an illustrative embodiment.

DETAILED DESCRIPTION

(9) Exemplary embodiments of the disclosure will now be discussed in further detail with regard to systems and methods for implementing intelligent load scheduling in a storage system. For purposes of illustration, exemplary embodiments will be described herein with reference to exemplary network computing systems, such as distributed storage systems, which implement data processing systems and associated computers, servers, storage devices and other processing devices. It is to be appreciated, however, that these and other embodiments are not restricted to the particular illustrative system and device configurations shown. Accordingly, the term “network computing system” as used herein is intended to be broadly construed, so as to encompass, for example, processing systems comprising cloud computing and storage systems, as well as other types of processing systems comprising various combinations of physical and virtual processing resources. A network computing system may therefore comprise, for example, at least one data center or other cloud-based systems that include one or more cloud systems that host multiple tenants which share cloud resources. Numerous different types of enterprise computing and storage systems are also encompassed by the term “network computing system” as that term is broadly used herein.

(10) As used herein, “real-time” refers to output within strict time constraints. Real-time output can be understood to be instantaneous or on the order of milliseconds or microseconds. Real-time output can occur when the connections with a network are continuous and a user device receives messages without any significant time delay. Of course, it should be understood that depending on the particular temporal nature of the system in which an embodiment is implemented, other appropriate timescales that provide at least contemporaneous performance and output can be achieved.

(11) FIG. 1 schematically illustrates a network computing system comprising a storage system which implements an intelligent load scheduling system, according to an exemplary embodiment of the disclosure. The network computing system **100** comprises one or more host systems **110-1, 110-2, . . . 110-H** (collectively, host systems **110**), one or more management nodes **115**, a communications network **120**, and a data storage system **130**. The data storage system **130** comprises one or more storage nodes **140-1, 140-2, . . . , 140-N** (collectively, storage nodes **140**),

and may comprise a distributed storage system. As shown in FIG. 1, the storage node **140-1** comprises a storage control system **150**, and a storage array **160** comprising a plurality of storage devices **162-1**, . . . , **162-D** (collectively, storage devices **162**). The storage control system **150** comprises a storage data client **151**, storage data server **152**, and an intelligent load scheduling system **154**, the functions of which will be described in further detail below. In some embodiments, the other storage nodes **140-2** . . . **140-N** have the same or similar configuration as the storage node **140-1** shown in FIG. 1.

(12) In general, the management nodes **115** implement application programming interfaces (APIs) to enable manual, automated, and/or semi-automated configuration, management, provisioning, and monitoring of the data storage system **130** and the associated storage nodes **140**. In some embodiments, the management nodes **115** comprise stand-alone dedicated management server nodes, which may comprise physical and/or virtual server nodes.

(13) The host systems **110** comprise physical server nodes and/or virtual server nodes which host and execute applications that are configured to process data and execute tasks/workloads and perform computational work, either individually, or in a distributed manner, to thereby provide compute services to one or more users (the term “user” herein is intended to be broadly construed so as to encompass numerous arrangements of human, hardware, software or firmware entities, as well as combinations of such entities). In some embodiments, the host systems **110** comprise application servers, database servers, etc. In some embodiments, the host systems **110** include virtual nodes such as virtual machines and container systems. In some embodiments, the host systems **110** comprise a cluster of computing nodes of an enterprise computing system, a cloud-based computing system, or other types of computing systems or information processing systems comprising multiple computing nodes associated with respective users. The host systems **110** issue data access requests to the data storage system **130**, wherein the data access requests include (i) write requests to store data in the storage arrays **160** of the storage nodes **140** and (ii) read requests to access data that is stored in storage arrays **160** of the storage nodes **140**. Such data access requests are referred to herein as “input-output (**10**) operations” that are processed by a corresponding one of the storage nodes **140**. The term “input-output” as used herein refers to at least one of input and output. For example, IO operations may comprise write requests and/or read requests directed to logical addresses of a particular logical storage volume of a given one of the storage nodes **140**. These and other types of IO operations are also generally referred to herein as IO requests.

(14) The communications network **120** is configured to enable communication between the host systems **110** and the storage nodes **140**, and between the management nodes **115**, the host systems **110**, and the storage nodes **140**, as well as to enable peer-to-peer communication between the storage nodes **140** of the data storage system **130**. In this regard, while the communications network **120** is generically depicted in FIG. 1, it is to be understood that the communications network **120** may comprise any known communications network such as, a global computer network (e.g., the Internet), a wide area network (WAN), a local area network (LAN), an intranet, a satellite network, a telephone or cable network, a cellular network, a wireless network such as Wi-Fi or WiMAX, a storage fabric (e.g., IP-based or Fiber Channel storage fabric), or various portions or combinations of these and other types of networks. In this regard, the term “network” as used herein is therefore intended to be broadly construed so as to encompass a wide variety of different network arrangements, including combinations of multiple networks possibly of different types, which enable communication using, e.g., Transfer Control Protocol/Internet Protocol (TCP/IP) or other communication protocols such as Fibre Channel (FC), FC over Ethernet (FCoE), Internet Small Computer System Interface (iSCSI), Peripheral Component Interconnect express (PCIe), InfiniBand, Gigabit Ethernet, etc., to implement input/output (I/O) channels and support storage network connectivity. Numerous alternative networking arrangements are possible in a given embodiment, as will be appreciated by those skilled in the art.

(15) The data storage system **130** may comprise any type of data storage system, or combination of data storage systems, including, but not limited to, a storage area network (SAN) system, a network-attached storage (NAS) system, a direct-attached storage (DAS) system, dynamic scale-out data storage systems, or other types of distributed data storage systems comprising software-defined storage, clustered or distributed virtual and/or physical infrastructure. The term “data storage system” as used herein should be broadly construed and not viewed as being limited to storage systems of any particular type or types. It is to be noted that each storage node **140** and its associated storage array **160** is an example of what is more generally referred to herein as a “storage system” or a “storage array.” The storage nodes **140** can be physical nodes, virtual nodes, and a combination of physical and virtual nodes.

(16) In some embodiments, the storage nodes **140** comprise storage server nodes (e.g., server node **800**, shown in FIG. **8**) which comprise processors and system memory, and host virtual machines and/or containers, although numerous other configurations are possible. In some embodiments, one or more of the storage nodes **140** can additionally implement functionality of a compute node, and vice-versa, wherein a compute node is configured to process data and execute tasks/workloads and perform computational work, either individually, or in a distributed manner, to thereby provide compute services such as execution of one or more applications on behalf of one or more users. In this regard, the term “storage node” as used herein is therefore intended to be broadly construed, and a storage system in some embodiments can be implemented using a combination of storage nodes and compute nodes.

(17) In some embodiments, each storage node **140** comprises a server node that is implemented on, e.g., a physical server machine or storage appliance comprising hardware processors, system memory, and other hardware resources that execute software and firmware to implement the functionalities and data management services of the storage node **140** and the storage control system **150**, as discussed herein. More specifically, in some embodiments, each storage node **140** comprises a plurality of storage control processors which execute a lightweight operating system (e.g., a customized lightweight Linux kernel) and functional software (e.g., software-defined storage software) to implement various functions of the storage node **140** and the storage control system **150**, wherein such functions include, but are not limited to, (i) managing and executing data access requests issued by the host systems **110**, (ii) performing various data management and storage services, and (iii) controlling network communication and connectivity with the host systems **110** and between the storage nodes **140** within the data storage system **130**, etc.

(18) In a distributed storage environment, the storage control systems **150** of the storage nodes **140** are configured to communicate in a cooperative manner to perform functions such as e.g., processing data access requests received from the host systems **110**, aggregating/pooling the storage capacity of storage arrays **160** of the storage nodes **140**, performing functions such as inline data compression/decompression, data deduplication, thin provisioning, and data protection functions such as data replication, snapshot, and data protection and resiliency schemes based on data striping and/or parity (e.g., erasure coding, RAID, etc.), and other types of data management functions, depending on the system configuration.

(19) The storage devices **162** comprise one or more of various types of storage devices such as hard-disk drives (HDDs), solid-state drives (SSDs), flash memory cards, or other types of non-volatile memory (NVM) devices including, but not limited to, non-volatile random-access memory (NVRAM), phase-change RAM (PC-RAM), magnetic RAM (MRAM), etc. In some embodiments, the storage devices **162** comprise flash memory devices such as NAND flash memory, NOR flash memory, etc. The NAND flash memory can include single-level cell (SLC) devices, multi-level cell (MLC) devices, triple-level cell (TLC) devices, or quad-level cell (QLC) devices. These and various combinations of multiple different types of storage devices may be implemented in the data storage system **130**. In this regard, the term “storage device” as used herein should be broadly construed to encompass all types of persistent storage media including hybrid drives.

(20) In some embodiments, the storage array **160** of a given storage node **140** comprises DAS resources (internal and/or external), wherein the storage control system **150** of the given storage node **140** is configured to directly access the storage array **160** of the given storage node **140**. In some embodiments, the data storage system **130** comprises a disaggregated data storage system in which storage data processing is separate from data storage. More specifically, in an exemplary embodiment of a disaggregated storage system, the storage control systems **150** comprise storage control nodes, and the storage arrays **160** comprise storage nodes, which are separate from the storage control nodes. In such a configuration, the storage control systems **150** are configured to handle the processing of data associated with data access requests (i.e., IO read and write requests), and the storage arrays **160** are configured to handle writing/reading data to/from respective storage devices **162**. In a disaggregated architecture, each storage control system **150** would be configured to directly access data stored in each storage array **160** in the data storage system **130**. The disaggregated storage system architecture essentially separates the storage control compute layers (e.g., storage control systems **150**) from the data storage layers (e.g., storage arrays **160**).

(21) In a disaggregated data storage system, each storage array **160** is implemented as, e.g., an external DAS device, wherein each storage control system **150** of each storage node **140-1**, **140-2**, . . . , **140-N** is connected to each storage array **160** using any suitable interface protocol such as Small Computer Systems Interface (SCSI), Fibre Channel (FC), etc. In other embodiments, the storage control systems **150** of each storage node **140-1**, **140-2**, . . . , **140-N** can be network-connected to each of the storage arrays **160** (via a high-performance network fabric) using any suitable network configuration and network interface protocol such as Ethernet, FC, Internet Small Computer Systems Interface (iSCSI), InfiniBand, etc. For example, in some embodiments, the storage control systems **150** and the storage arrays **160** are interconnected in a full-mesh network, wherein back-end interconnectivity is achieved using, e.g., a redundant high-speed storage fabric, wherein the storage control systems **150** can utilize remote procedure calls (RPC) for control messages and remote direct memory access (RDMA) for accessing data blocks.

(22) In some embodiments, the storage data servers **152** of the storage nodes **140** are configured to consolidate the capacity of the storage arrays **160** (e.g., HDDs, SSDs, PCIe or NVMe flash cards, etc.) of the storage nodes **140** into storage pools from which logical volumes are allocated, wherein the logical volumes (e.g., a block unit of storage management) are identified by, e.g., logical unit numbers (LUNs). More specifically, the storage data servers **152** of the storage nodes **140** are configured to create and manage storage pools (e.g., virtual pools of block storage) by aggregating storage capacity of the storage arrays **160** of the storage nodes **140** and dividing a given storage pool into one or more volumes, wherein the volumes are exposed to the host systems **110** as block devices. For example, a virtual block device can correspond to a volume of a storage pool. Each virtual block device comprises any number of actual physical storage devices, wherein each block device is preferably homogenous in terms of the type of storage devices that make up the block device (e.g., a block device can include only HDD devices or SSD devices, etc.).

(23) In some embodiments, each host system **110** comprises a storage data client (SDC) which executes on the host system and which consumes the block storage exposed by the storage data servers **152**. In particular, an SDC comprises a lightweight block device driver that is deployed on a given host system **110** to expose shared block volumes to the given host system **110**. The SDC exposes the storage volumes as block devices to each application (e.g., virtual machine, container, etc.) that execute on the same server (e.g., host system **110**) on which the SDC is installed. The SDC of a given host system **110** exposes block devices representing the virtual storage volumes that are currently mapped to the given host system **110**. The SDC for a given host system **110** serves as a block driver for the host system **110**, wherein the SDC intercepts **10** requests, and utilizes the intercepted **10** request to access the block storage that is managed by the storage data servers **152**. The SDC provides the operating system or hypervisor (which runs the SDC) access to the logical block devices (e.g., volumes). Each SDC has knowledge of which storage data servers

152 hold (e.g., own) their block data, so multipathing can be accomplished natively through the SDCs. Alternatively, or in addition to the SDCs on the host systems **110**, as shown in FIG. **1**, each of the storage nodes **140** comprises a storage data client (SDC) **151**, which performs the same or substantially the same functions as the SDCs on the host systems **110**.

(24) As noted above, the management nodes **115** in FIG. **1** implement a management layer which manages and configures the network computing system **100**. In some embodiments, the management nodes **115** comprise a tightly-coupled cluster of manager nodes that are configured to supervise the operations of the storage cluster and manage storage cluster configurations. For example, management nodes **115** include metadata manager (MDM) modules that operate outside of the data path and provide the relevant information to the SDCs on the host systems **110** and/or SDCs **151**, and to the storage data servers **152** to allow such components to control data path operations. The MDM modules are configured to manage the mapping of the SDCs on the host systems **110** and/or SDCs **151** to the storage data servers **152** of the storage nodes **140**. The MDM modules manage various types of metadata that are required to perform various management operations in the storage environment such as, e.g., managing configuration changes, managing the SDCs on the host systems **110** and/or SDCs **151** and storage data servers **152**, maintaining and updating device mappings, maintaining management metadata for controlling data protection operations such as snapshots, replication, RAID configurations, etc., managing system capacity including device allocations and/or release of capacity, performing operation for recovery from errors and failures, and system rebuild tasks including rebalancing, etc.

(25) Workloads from host systems **110** (e.g., IO operations) are not the only workloads the storage nodes **140** are executing. The storage control system **150**, more particularly, the intelligent load scheduling system **154**, performs a workload balancing process that accounts for total storage node workload comprising foreground processes (i.e., IO operations) and background processes (e.g., data checksum validation, defragmentation, scrubbing, data compression/decompression, data deduplication, thin provisioning, data protection functions, etc.). Other workload service level vectors, such as, for example, reliability requirements, up-time, performance, etc. may also be addressed by background processes. Workload balancing may be performed on an individual storage node **140** or between multiple storage nodes **140** such as in a distributed storage environment.

(26) Processing resources vary between storage nodes **140**. For example, storage nodes **140** with weaker CPU cores than other storage nodes **140** will experience some workloads as heavier loads than the other storage nodes **140** with stronger CPU cores. In other examples, storage nodes **140** having graphics processing units (GPUs) and/or more memory may be able to better process certain workloads than storage nodes **140** without GPUs and/or with less memory.

(27) To account for the variable hardware and other resource capabilities between storage nodes **140**, the embodiments advantageously provide techniques for controlling loads on a storage node **140**. In particular, the embodiments provide an intelligent load scheduling system **154** which is configured to balance the loads from foreground processes (e.g., IO operations) and background processes (e.g., scrubbing, defragmentation, etc.) in order not to impact the foreground processes. Foreground processes are typically created by user (e.g., customer or client) applications and may have higher priority than background processes. Illustrative embodiments are advantageously configured to perform real-time analysis of storage node loads, and schedule background processes to execute later than foreground processes when necessary. The intelligent load scheduling system **154** uses machine learning techniques to determine when deferral of background processes is needed, and when to reschedule the background processes. In more detail, the embodiments utilize Principal Component Analysis (PCA) techniques to generate a function which enables storage nodes **140** to implement load scheduling in real-time responsive to incoming workloads from host systems **110**. Exemplary systems and methods for performing intelligent load scheduling according to embodiments of the disclosure will now be discussed in further detail in conjunction with FIGS.

2-7.

(28) It is to be understood that although the embodiments are discussed in terms of storage nodes **140**, the embodiments are not necessarily limited thereto, and are applicable to other types of servers where load scheduling as described herein may need to be implemented.

(29) FIG. 2 schematically illustrates a storage node which implements an intelligent load scheduling system, according to an exemplary embodiment of the disclosure. In some embodiments, FIG. 2 schematically illustrates an exemplary architecture of the storage nodes **140** of the data storage system **130** of FIG. 1. As shown in FIG. 2, the storage node **200** comprises a storage control system **210** which implements a storage data server **220**, a storage data client **224**, a data management services module **230**, a cache management module **240**, and an intelligent load scheduling system **250**. The storage data server **220** comprises a storage virtualization management module **222**. The intelligent load scheduling system **250** comprises various modules including, but not limited to, a metrics collection module **251**, a PCA module **252**, a load prediction module **253**, a training module **254**, and a background (BG) process scheduling module **255**, the functions of which will be described in further detail below.

(30) The storage node **200** further comprises a storage array **260** and memory **270**. The storage array **260** comprises a plurality of storage devices, wherein each storage device comprises physical storage addresses that can be mapped to logical addresses of one or more volumes **262** (or portions of volumes) that are configured within the storage array **260**. In some embodiments, the volume(s) **262** have logical addresses that are partitioned into extents, and each storage device comprises physical addresses that are mappable to logical addresses such that each storage device comprises extents corresponding to a portion of the logical address space. In some embodiments, the memory **270** comprises volatile random-access memory (RAM) such as dynamic RAM (DRAM), synchronous DRAM (SDRAM), etc. The memory **270** is configured to store, for example, configuration files that are utilized to support intelligent load scheduling operations as discussed herein.

(31) The storage data server **220** implements functions as discussed above in connection with the storage data server **152**, such as processing **10** write and read requests received from host systems to write/read data to/from the storage array **260**. The storage virtualization management module **222** implements any suitable logical volume management system which is configured to create and manage the volume(s) **262** by aggregating the capacity of the storage array **260** into one or more virtual storage pools that are thin-provisioned for maximum capacity, and logically dividing each storage pool into one or more storage volumes that are exposed as block devices (e.g., LUNs) to the applications or host systems **110** (FIG. 1) which consume the data. The storage data client **224** implements functions as discussed above in connection with the SDCs **151** such as, for example, exposing shared block volumes to a given host system. The cache management module **240** provisions and manages one or more in-RAM caches (e.g., read cache) in the memory **270**.

(32) The data management services module **230** implements one or more types of data management services including, but not limited to, inline data compression/decompression, thin provisioning, and data protection functions such as data replication, data backup, data snapshot, and data protection and resiliency schemes based on data striping and/or parity (e.g., erasure coding, RAID, etc.), and other types of data management functions, depending on the system configuration. In embodiments where the storage data server **220** abstracts the physical media (e.g., storage array **260**) and presents logical (virtualized) addresses to users in the form of LUNs, the storage data server **220** generates and manages metadata to provide mapping between logical addresses and physical addresses (e.g., a logical-to-physical mapping directory) which is stored on disk, wherein at any given time, a plurality of logical-to-physical pages of the logical-to-physical mapping directory are swapped into the memory **270** and stored in a read cache during runtime (e.g., the cached logical-to-physical pages). In addition, the storage control system **210** generates and manages metadata which is utilized for managing snapshots, managing remote replication,

managing deduplication, managing data compression, and maintaining resiliency related metadata (e.g., RAID), etc.

(33) The various modules of the intelligent load scheduling system **250** implement functions to enable the storage control system **210** to perform load scheduling operations on the storage nodes **200** (or **140**). The metrics collection module **251** for each storage node **200** collects a set of measured metrics that correlate to one or more workloads (e.g., IO operations) received from one or more host devices of a host system. The metrics comprise, for example, a write size rate, a read size rate and a latency rate for a storage node **200**. In some cases, the metrics are codependent. As a result, simple arithmetic techniques (e.g., mean) are not used to infer a current workload of a storage node **200**. Depending on whether the intelligent load scheduling system **250** is in a training phase, or a test phase, the collected metrics may be, for example, training or test metrics.

(34) The PCA module **252**, in conjunction with the training module **254**, processes a plurality of training metrics to reduce one or more dimensions of and remove dependencies from the metrics. During training, the metrics collection module **251** collects the training metrics for one or more workloads and identifies a current load level (e.g., percent (100%, 75%, 50%, etc.) loaded) on a storage node **200** or plurality of storage nodes **200** that is associated with the training metrics. Using PCA techniques, the PCA and training modules **252** and **254** process the training metrics and the current load level to allow PCA to produce respective weights for the supplied training metrics. In a non-limiting illustrative example, a metric X has a weight of 70% (e.g., 70% of the metric X contributes to the load), while a metric Y weight has a weight of 5% (e.g., 5% of metric Y contributes to the load). The PCA module **252** arranges the weights in a formula (weight*test metric value), which is used by the load prediction module **253** to predict an overall load on one or more storage nodes **200** when test metrics are collected. The load prediction module **253** is configured to compute load on a per storage node or system-wide basis.

(35) Using the weights for respective ones of the plurality of training metrics generated during training, the load prediction module **253** processes a plurality of test metrics to predict a per storage node or system-wide load corresponding to the test metrics. Based on the predicted load, the BG process scheduling module **255** automatically determines the allowed background processes to be executed, and the background processes to be postponed to a later time. The intelligent load scheduling system collects test metrics, predicts a load and schedules background processes in real-time (e.g., within microseconds or milliseconds).

(36) The PCA module **252** accounts for all metrics associated with workloads from host systems (e.g., host systems **110**) and provides an aggregated result that gives each metric its own weight in the decision when the next workload (e.g., background operation) should run. For example, there is a load correlation/dependency between a device's (e.g., storage node's) read rate and the sizes of its read commands. Another example for parameter dependency is defragmentation IOs, which are added to foreground IO operations. As a result, a total storage node load must account for defragmentation IOs by subtracting them from foreground IO operations. Without the techniques of the embodiments (e.g., PCA or other equivalent machine learning methodology) to reduce dependent variables, a load calculation algorithm may include a bias and, in some cases, require more resources to process the load calculation than to perform the workloads.

(37) After the PCA module **252** processes the training metrics, the results of the multiple linear regression (e.g., a set of respective weights respectively corresponding to the metrics) are used by the load prediction module **253** to predict the per storage node or system-wide load when scheduling operations (e.g., background processes) based on test metrics corresponding to incoming workloads. For example, when considering starting a new background operation, the metrics collection module **251** will collect a set of test metrics and the load prediction module **253** uses the weights generated during training to compute a current load, which is used by the BG process scheduling module **255** to determine whether the new background operation may commence. For example, the BG process scheduling module **255** compares the predicted load to a

designated load threshold. If the predicted load is below the designated load threshold, the background operation can start. If the predicted load is at or above the designated load threshold, then the background operation is deferred to a later time.

(38) According to illustrative embodiments, the constants (e.g., weights) are generated through a statistical analysis on benchmark measurements used by the PCA and training modules **252** and **254** during training. As noted herein above, PCA techniques are used by the PCA module **252** to reduce dependencies so that the resulting adaptive function for real-time workloads analysis will have better accuracy and consume less processing resources. In other words, the weights produced by the PCA techniques, which are used with the test metrics by the load prediction module **253** to determine current load, produce an accurate per storage node or system-wide load using simpler computations than those that would be required without use of the PCA or other reduction techniques of the illustrative embodiments.

(39) The PCA techniques used by the PCA module **252** reduce dimensionality and remove feature dependence. The principal components are used as vectors. If the vectors are orthogonal, they are independent, and thus sorted by importance to the overall storage node load.

(40) For example, the amount of data described by $\text{PRINCIPAL_COMPONENT}[i] > \text{the amount of data described by PRINCIPAL_COMPONENT}[i+1]$ where $0 \leq i < n$. In illustrative embodiments, the number of principal components for an output is configurable. For example, assuming there are 4 metrics and at least an 85% accuracy is required. If the amount of data described by the principal components is as follows: [70%, 13%, 11%, 6%] (where all the percentages add up to 100%), then 3 principal components are needed to add up to at least 85% (e.g., 70%, 13% and 11%; 70%, 13% and 6%; or 70%, 11% and 6%). In this case, the dimensionality is reduced from **4** to **3**.

(41) If dependent features are the measured metrics (e.g., device write/read size rate, etc.), the dimension is the number of measured metrics and the dimension to reduce to is the number of principal components. After running the training data through the PCA model, a new set of metrics that are in the PCA space is generated, which can be used in multiple regression to find a linear fit.

(42) FIG. 3 depicts a plot **300** of an output of multiple linear regression on metrics processed using PCA, where two principal components were requested. The lighter (gray) portions correspond to the actual measurements of the variable being predicted (e.g., System/Device Load) and the darker (black) portions correspond to the predicted load provided by the load prediction module **253** given by multiple regression. The units for load in this example are 10 operations per second (IOPS) in megabytes per second (MBps). As can be seen in FIG. 3, the x-axis corresponds to Principal Component 0, the y-axis corresponds to Principal Component 1 and z-axis corresponds to IOPS primary writes (Standardized). All of the black points are in the 2-dimensional plane.

(43) FIG. 4 depicts a table **400** of an output of multiple linear regression on metrics processed using PCA. Based on the data in the table **400**, two principal components explain 79% of the data for which prediction is being attempted. A coefficient of determination (R-squared) is 0.866, which indicates a good fit. As can be understood, the coefficient of determination corresponds to the proportion of the variation in the dependent variable that is predictable from the independent variable(s), and varies between 0 and 1, where a number closer to 1 indicates a good fit. For example, in percentage terms, 0.866 means that the model explains 86.6% of the variation in the dependent variable. The principal components are a set of coefficients that define a linear combination of the measured metrics.

(44) FIG. 5 depicts a plot **500** of an output of multiple linear regression on metrics that are different from those used in connection with FIGS. 3 and 4 and processed using PCA, where two principal components were requested. Similar to FIG. 3, the lighter (gray) portions correspond to the actual measurements of the variable being predicted (e.g., System/Device Load) and the darker (black) portions correspond to the predicted load provided by the load prediction module **253** given by multiple regression. Similar to FIG. 3, the x-axis corresponds to Principal Component 0, the y-axis corresponds to Principal Component 1 and z-axis corresponds to IOPS primary writes

(Standardized). FIG. 6 depicts an alternative plot **600** of the output of multiple linear regression on the metrics of FIG. 5 in an illustrative embodiment. Similar to FIGS. 3 and 5, the x-axis corresponds to Principal Component 0, the y-axis corresponds to Principal Component 1 and z-axis corresponds to IOPS primary writes (Standardized).

(45) According to illustrative embodiments, for each new measurement, the function is refining its coefficients, making it adaptive to the storage node **200** (or **140**) it runs on. For example, in a training loop, over relatively long time intervals (e.g., longer than the interval for running the inference loop), collected training metrics and the corresponding per storage node or system-wide load are input to the PCA module **252**. The PCA module **252** runs a multiple linear regression to update the linear inference function coefficients (e.g., metric weights), which are updated in a configuration file stored in the memory **270**. In the inference loop, over relatively short intervals (e.g., 2-3 milliseconds) the metrics collection module **251** collects a plurality of test metrics corresponding to one or more workloads. The metrics are applied to the linear inference function by the load prediction module **253** to predict per storage node or system-wide load, and the BG process scheduling module **255** uses the predicted load to determine scheduling of background processes based on whether predicted load exceeds a designated threshold load.

(46) Using more parameters increases the accuracy of the linear inference function. However, repeatedly collecting and analyzing metrics uses system resources. As a result, in some embodiments, to limit the load associated with metrics collection, a sliding window of a designated time (e.g., an X second sliding window) will be used to aggregate the parameters, where X is an integer greater than 0, and is configurable. In another alternative to limit the load associated with metrics collection, the metrics for each storage node are collected at designated intervals corresponding to a number of IO operations (e.g., every Y IO operations), where Y is an integer greater than 0 (e.g., Y=1000), and is configurable. With this approach, devices without 10 operations will not add load.

(47) An exemplary method for implementing intelligent load scheduling will now be discussed in conjunction with the flow diagram of FIG. 7. In some embodiments, the process flow shown in FIG. 7 illustrates exemplary modes of operation of the intelligent load scheduling system **250** of FIG. 2. Referring to FIG. 7, at block **700**, the intelligent load scheduling system **250** collects a plurality of metrics corresponding to one or more workloads. The one or more workloads correspond to one or more IO operations from one or more host devices. The collecting may be performed at designated intervals corresponding to a number of the IO operations.

(48) At block **702**, the intelligent load scheduling system **250** processes the plurality of metrics using a machine learning model to predict a load on at least one server (e.g., storage node). The machine learning model (e.g., PCA model) predicts the load on the at least one server based, at least in part, on weights for respective ones of the plurality of metrics generated during training of the machine learning model. The training of the machine learning model comprises collecting a plurality of training metrics, identifying a load on the at least one server corresponding to the plurality of training metrics, and processing the plurality of training metrics and the identified load by the machine learning model to generate the weights. The processing comprises running the plurality of training metrics and the identified load through a multiple linear regression. The training of the machine learning model reduces one or more dimensions of the plurality of training metrics, and/or eliminates one or more dependencies corresponding to the plurality of training metrics.

(49) At block **704**, the intelligent load scheduling system **250** schedules execution by the at least one server of one or more processes in addition to the one or more workloads based, at least in part, on the load predicted by the machine learning model. The one or more processes correspond to one or more background processes such as, but not necessarily limited to, data checksum validation and defragmentation. The scheduling of the execution by the at least one server of the one or more processes in addition to the one or more workloads comprises deferring the one or more processes.

(50) It is to be appreciated that the FIG. 7 process and other features and functionality described above can be adapted for use with other types of systems configured to perform load scheduling in storage systems.

(51) The particular processing operations and other system functionality described in conjunction with the flow diagram of FIG. 7 is therefore presented by way of illustrative example only, and should not be construed as limiting the scope of the disclosure in any way. Alternative embodiments can use other types of processing operations. For example, the ordering of the process steps may be varied in other embodiments, or certain steps may be performed at least in part concurrently with one another rather than serially. Also, one or more of the process steps may be repeated periodically, or multiple instances of the process can be performed in parallel with one another.

(52) Functionality such as that described in conjunction with the flow diagram of FIG. 7 can be implemented at least in part in the form of one or more software programs stored in memory and executed by a processor of a processing device such as a computer or server. As will be described below, a memory or other storage device having executable program code of one or more software programs embodied therein is an example of what is more generally referred to herein as a “processor-readable storage medium” or “processor-readable storage media”.

(53) Illustrative embodiments provide technical solutions that use machine learning techniques (e.g., PCA) to generate an adaptive function for inferring real-time loads, and use the predicted loads when determining whether to execute background processes. Unlike conventional approaches, the embodiments provide techniques which reduce dimensions and remove dependencies of input metrics to allow for rapid (e.g., real-time) determinations of device and system loads without unnecessary use of compute resources.

(54) The embodiments address concerns of potential bias and excessive use of compute resources by generating an adaptive linear function that gives each metric its own weight for load determination, and by automatically applying predicted loads to an algorithm for determining when background processes should execute. As a result, the embodiments provide a methodology for automated adaptive throttling of background processes that is independent of the system topology and resource variety.

(55) It is to be appreciated that the particular advantages described above and elsewhere herein are associated with particular illustrative embodiments and need not be present in other embodiments. Also, the particular types of information processing systems and/or features and functionality as illustrated in the drawings and described above are exemplary only, and numerous other arrangements may be used in other embodiments.

(56) FIG. 8 schematically illustrates a framework of a server for hosting a storage control system which implements an intelligent load scheduling system, according to an exemplary embodiment of the disclosure. The server node **800** comprises processors **802**, storage interface circuitry **804**, network interface circuitry **806**, virtualization resources **808**, system memory **810**, and storage resources **816**. The system memory **810** comprises volatile memory **812** and non-volatile memory **814**. The processors **802** comprise one or more types of hardware processors that are configured to process program instructions and data to execute a native operating system (OS) and applications that run on the server node **800**.

(57) For example, the processors **802** may comprise one or more CPUs, microprocessors, microcontrollers, application specific integrated circuits (ASICs), field programmable gate arrays (FPGAs), and other types of processors, as well as portions or combinations of such processors. The term “processor” as used herein is intended to be broadly construed so as to include any type of processor that performs processing functions based on software, hardware, firmware, etc. For example, a “processor” is broadly construed so as to encompass all types of hardware processors including, for example, (i) general purpose processors which comprise “performance cores” (e.g., low latency cores), and (ii) workload-optimized processors, which comprise any possible

combination of multiple “throughput cores” and/or multiple hardware-based accelerators.

Examples of workload-optimized processors include, for example, graphics processing units (GPUs), digital signal processors (DSPs), system-on-chip (SoC), tensor processing units (TPUs), image processing units (IPUs), deep learning accelerators (DLAs), artificial intelligence (AI) accelerators, and other types of specialized processors or coprocessors that are configured to execute one or more fixed functions.

(58) The storage interface circuitry **804** enables the processors **802** to interface and communicate with the system memory **810**, the storage resources **816**, and other local storage and off-infrastructure storage media, using one or more standard communication and/or storage control protocols to read data from or write data to volatile and non-volatile memory/storage devices. Such protocols include, but are not limited to, NVMe, PCIe, PATA, SATA, SAS, Fibre Channel, etc. The network interface circuitry **806** enables the server node **800** to interface and communicate with a network and other system components. The network interface circuitry **806** comprises network controllers such as network cards and resources (e.g., network interface controllers (NICs) (e.g., SmartNICs, RDMA-enabled NICs), Host Bus Adapter (HBA) cards, Host Channel Adapter (HCA) cards, IO adaptors, converged Ethernet adaptors, etc.) to support communication protocols and interfaces including, but not limited to, PCIe, DMA and RDMA data transfer protocols, etc.

(59) The virtualization resources **808** can be instantiated to execute one or more services or functions which are hosted by the server node **800**. For example, the virtualization resources **808** can be configured to implement the various modules and functionalities of a host connectivity management system as discussed herein. In some embodiments, the virtualization resources **808** comprise virtual machines that are implemented using a hypervisor platform which executes on the server node **800**, wherein one or more virtual machines can be instantiated to execute functions of the server node **800**. As is known in the art, virtual machines are logical processing elements that may be instantiated on one or more physical processing elements (e.g., servers, computers, or other processing devices). That is, a “virtual machine” generally refers to a software implementation of a machine (i.e., a computer) that executes programs in a manner similar to that of a physical machine. Thus, different virtual machines can run different operating systems and multiple applications on the same physical computer.

(60) A hypervisor is an example of what is more generally referred to as “virtualization infrastructure.” The hypervisor runs on physical infrastructure, e.g., CPUs and/or storage devices, of the server node **800**, and emulates the CPUs, memory, hard disk, network and other hardware resources of the host system, enabling multiple virtual machines to share the resources. The hypervisor can emulate multiple virtual hardware platforms that are isolated from each other, allowing virtual machines to run, e.g., Linux and Windows Server operating systems on the same underlying physical host. The underlying physical infrastructure may comprise one or more commercially available distributed processing platforms which are suitable for the target application.

(61) In other embodiments, the virtualization resources **808** comprise containers such as Docker containers or other types of Linux containers (LXCs). As is known in the art, in a container-based application framework, each application container comprises a separate application and associated dependencies and other components to provide a complete filesystem, but shares the kernel functions of a host operating system with the other application containers. Each application container executes as an isolated process in user space of a host operating system. In particular, a container system utilizes an underlying operating system that provides the basic services to all containerized applications using virtual-memory support for isolation. One or more containers can be instantiated to execute one or more applications or functions of the server node **800** as well execute one or more of the various modules and functionalities of a storage control node and intelligent load scheduling system as discussed herein. In yet another embodiment, containers may be used in combination with other virtualization infrastructure such as virtual machines

implemented using a hypervisor, wherein Docker containers or other types of LXC's are configured to run on virtual machines in a multi-tenant environment.

(62) In some embodiments, the constituent components and modules of the storage control systems and intelligent load scheduling systems are implemented using program code that is loaded into the system memory **810** (e.g., volatile memory **812**), and executed by the processors **802** to perform respective functions as described herein. In this regard, the system memory **810**, the storage resources **816**, and other memory or storage resources as described herein, which have program code and data tangibly embodied thereon, are examples of what is more generally referred to herein as “processor-readable storage media” that store executable program code of one or more software programs. Articles of manufacture comprising such processor-readable storage media are considered embodiments of the disclosure. An article of manufacture may comprise, for example, a storage device such as a storage disk, a storage array or an integrated circuit containing memory. The term “article of manufacture” as used herein should be understood to exclude transitory, propagating signals.

(63) The system memory **810** comprises various types of memory such as volatile RAM, NVRAM, or other types of memory, in any combination. The volatile memory **812** may be a dynamic random-access memory (DRAM) (e.g., DRAM DIMM (Dual In-line Memory Module)), or other forms of volatile RAM. The non-volatile memory **814** may comprise one or more of NAND Flash storage devices, SSD devices, or other types of next generation non-volatile memory (NGNVM) devices. The system memory **810** can be implemented using a hierarchical memory tier structure wherein the volatile memory **812** is configured as the highest-level memory tier, and the non-volatile memory **814** (and other additional non-volatile memory devices which comprise storage-class memory) is configured as a lower level memory tier which is utilized as a high-speed load/store non-volatile memory device on a processor memory bus (i.e., data is accessed with loads and stores, instead of with **10** reads and writes). The term “memory” or “system memory” as used herein refers to volatile and/or non-volatile memory which is utilized to store application program instructions that are read and processed by the processors **802** to execute a native operating system and one or more applications or processes hosted by the server node **800**, and to temporarily store data that is utilized and/or generated by the native OS and application programs and processes running on the server node **800**. The storage resources **816** can include one or more HDDs, SSD storage devices, etc.

(64) It is to be understood that the above-described embodiments of the disclosure are presented for purposes of illustration only. Many variations may be made in the particular arrangements shown. For example, although described in the context of particular system and device configurations, the techniques are applicable to a wide variety of other types of information processing systems, computing systems, data storage systems, processing devices and distributed virtual infrastructure arrangements. In addition, any simplifying assumptions made above in the course of describing the illustrative embodiments should also be viewed as exemplary rather than as requirements or limitations of such embodiments. Numerous other alternative embodiments within the scope of the appended claims will be readily apparent to those skilled in the art.

Claims

1. A method, comprising: collecting a plurality of metrics corresponding to one or more workloads at designated intervals corresponding to a number of the one or more workloads; processing the plurality of metrics using a machine learning model to predict a load on at least one server, wherein the machine learning model predicts the load on the at least one server based, at least in part, on weights for respective ones of the plurality of metrics generated during training of the machine learning model, wherein the training of the machine learning model comprises: collecting a plurality of training metrics; identifying the load on the at least one server corresponding to the

- plurality of training metrics; and processing the plurality of training metrics and the identified load by the machine learning model to generate the weights; and scheduling execution by the at least one server of one or more processes in addition to the one or more workloads based, at least in part, on the load predicted by the machine learning model; wherein the method is performed by at least one processing device comprising a processor coupled to a memory.
2. The method of claim 1, wherein the one or more workloads correspond to one or more input-output (IO) operations from one or more host devices.
 3. The method of claim 2, wherein the one or more processes correspond to one or more background processes.
 4. The method of claim 3, wherein the one or more background processes comprise at least one of data checksum validation and defragmentation.
 5. The method of claim 1, wherein the processing comprises running the plurality of training metrics and the identified load through a multiple linear regression.
 6. The method of claim 1, wherein the training of the machine learning model reduces one or more dimensions of the plurality of training metrics.
 7. The method of claim 1, wherein the training of the machine learning model eliminates one or more dependencies corresponding to the plurality of training metrics.
 8. The method of claim 1, wherein the machine learning model comprises a Principal Component Analysis (PCA) model.
 9. The method of claim 1, wherein the plurality of metrics comprise at least one of a write size rate, a read size rate and a latency rate of the at least one server.
 10. The method of claim 1, wherein the scheduling of the execution by the at least one server of the one or more processes in addition to the one or more workloads comprises deferring the one or more processes.
 11. An article of manufacture comprising a non-transitory processor-readable storage medium having stored therein program code of one or more software programs, wherein the program code is executable by one or more processors to implement a method which comprises: collecting a plurality of metrics corresponding to one or more workloads at designated intervals corresponding to a number of the one or more workloads; processing the plurality of metrics using a machine learning model to predict a load on at least one server, wherein the machine learning model predicts the load on the at least one server based, at least in part, on weights for respective ones of the plurality of metrics generated during training of the machine learning model, wherein the training of the machine learning model comprises: collecting a plurality of training metrics; identifying the load on the at least one server corresponding to the plurality of training metrics; and processing the plurality of training metrics and the identified load by the machine learning model to generate the weights; and scheduling execution by the at least one server of one or more processes in addition to the one or more workloads based, at least in part, on the load predicted by the machine learning model.
 12. The article of manufacture of claim 11, wherein the training of the machine learning model reduces one or more dimensions of the plurality of training metrics.
 13. The article of manufacture of claim 11, wherein the training of the machine learning model eliminates one or more dependencies corresponding to the plurality of training metrics.
 14. The article of manufacture of claim 11, wherein the processing comprises running the plurality of training metrics and the identified load through a multiple linear regression.
 15. The article of manufacture of claim 11, wherein the training of the machine learning model reduces one or more dimensions of the plurality of training metrics.
 16. An apparatus, comprising: at least one processing device comprising a processor coupled to a memory; said at least one processing device being configured to: collect a plurality of metrics corresponding to one or more workloads at designated intervals corresponding to a number of the one or more workloads; process the plurality of metrics using a machine learning model to predict a

load on at least one server, wherein the machine learning model predicts the load on the at least one server based, at least in part, on weights for respective ones of the plurality of metrics generated during training of the machine learning model, wherein the training of the machine learning model comprises: collecting a plurality of training metrics; identifying the load on the at least one server corresponding to the plurality of training metrics; and processing the plurality of training metrics and the identified load by the machine learning model to generate the weights; and schedule execution by the at least one server of one or more processes in addition to the one or more workloads based, at least in part, on the load predicted by the machine learning model.

17. The apparatus of claim 9, wherein the training of the machine learning model reduces one or more dimensions of the plurality of training metrics.

18. The apparatus of claim 9, wherein the training of the machine learning model eliminates one or more dependencies corresponding to the plurality of training metrics.

19. The apparatus of claim 16, wherein the processing comprises running the plurality of training metrics and the identified load through a multiple linear regression.

20. The apparatus of claim 16, wherein the training of the machine learning model reduces one or more dimensions of the plurality of training metrics.
