

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386684
Kind Code	B1
Date of Patent	August 12, 2025
Inventor(s)	Hinrichs; Timothy L. et al.

System for authorizing API calls

Abstract

Some embodiments of the invention provide a system for defining, distributing and enforcing policies for authorizing API (Application Programming Interface) calls to applications executing on one or more sets of associated machines (e.g., virtual machines, containers, computers, etc.) in one or more datacenters. This system has servers that act as a logically centralized resource for defining and storing policies and parameters for evaluating these policies. The servers enforce these policies and distribute the policies and parameters to policy-enforcing local agents that execute near the applications that process the API calls. From an associated application, a local agent receives API-authorization requests to determine whether API calls received by the application are authorized. In response to such a request, the local agent uses one or more parameters associated with the API call to identify a policy stored in its local policy storage to evaluate whether the API call should be authorized. To evaluate this policy, the agent also retrieves one or more parameters from the local policy storage.

Inventors: Hinrichs; Timothy L. (Los Altos, CA), Koponen; Teemu (San Francisco, CA), Curtis; Andrew (San Mateo, CA), Sandall; Torin (San Francisco, CA), Florescu; Octavian (Kirkland, WA)

Applicant: Styra, Inc. (Redwood City, CA)

Family ID: 69779221

Assignee: STYRA, INC. (Redwood City, CA)

Appl. No.: 16/050119

Filed: July 31, 2018

Related U.S. Application Data

us-provisional-application US 62545458 20170814
us-provisional-application US 62540547 20170802

Publication Classification

Int. Cl.: **G06F9/54** (20060101); **G06F9/455** (20180101); **G06F16/11** (20190101); **G06F21/30** (20130101)

U.S. Cl.:

CPC **G06F9/54** (20130101); **G06F9/45533** (20130101); **G06F16/122** (20190101); **G06F21/30** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5974549	12/1998	Golan	N/A	N/A
6460141	12/2001	Olden	N/A	N/A
6985953	12/2005	Sandhu et al.	N/A	N/A
7096367	12/2005	Garg et al.	N/A	N/A
7124192	12/2005	High, Jr. et al.	N/A	N/A
7752661	12/2009	Hemsath et al.	N/A	N/A
7865931	12/2010	Stone	726/17	G06F 21/31
7913300	12/2010	Flank et al.	N/A	N/A
7937755	12/2010	Guruswamy	N/A	N/A
8266694	12/2011	Roy	N/A	N/A
8613070	12/2012	Borzycki et al.	N/A	N/A
8683560	12/2013	Brooker et al.	N/A	N/A
8782744	12/2013	Fuller et al.	N/A	N/A
8789138	12/2013	Reierson et al.	N/A	N/A
8850528	12/2013	Biljon et al.	N/A	N/A
9106661	12/2014	Stamos	N/A	N/A
9374417	12/2015	Greenfield et al.	N/A	N/A
9397990	12/2015	Taly et al.	N/A	N/A
9420002	12/2015	McGovern et al.	N/A	N/A
9521032	12/2015	Worsley	N/A	N/A
9530020	12/2015	Brandwine et al.	N/A	N/A
9578004	12/2016	Greenspan et al.	N/A	N/A
9648040	12/2016	Morkel et al.	N/A	N/A
9948681	12/2017	Kruse et al.	N/A	N/A
10021103	12/2017	Xu et al.	N/A	N/A
10122757	12/2017	Kruse et al.	N/A	N/A
10127393	12/2017	Ferraiolo et al.	N/A	N/A
10148493	12/2017	Ennis, Jr. et al.	N/A	N/A
10182129	12/2018	Peterson et al.	N/A	N/A
10257184	12/2018	Mehta et al.	N/A	N/A
10263995	12/2018	Kruse et al.	N/A	N/A

10339303	12/2018	Mehta et al.	N/A	N/A
10353726	12/2018	Duan	N/A	N/A
10454975	12/2018	Sharifi Mehr	N/A	N/A
10469314	12/2018	Ennis, Jr. et al.	N/A	N/A
10574699	12/2019	Baer et al.	N/A	N/A
10592302	12/2019	Hinrichs et al.	N/A	N/A
10592683	12/2019	Lim et al.	N/A	N/A
10715514	12/2019	Threlkeld	N/A	N/A
10719373	12/2019	Koponen et al.	N/A	N/A
10740470	12/2019	Ionescu et al.	N/A	N/A
10789220	12/2019	Mayer et al.	N/A	N/A
10984133	12/2020	Hinrichs et al.	N/A	N/A
10986131	12/2020	Kruse et al.	N/A	N/A
10990702	12/2020	Hinrichs et al.	N/A	N/A
11023292	12/2020	Hinrichs et al.	N/A	N/A
11080410	12/2020	Sandall et al.	N/A	N/A
11108827	12/2020	Beckman et al.	N/A	N/A
11108828	12/2020	Curtis et al.	N/A	N/A
11170099	12/2020	Sandall et al.	N/A	N/A
11228573	12/2021	Rangasamy et al.	N/A	N/A
11258824	12/2021	Hinrichs et al.	N/A	N/A
11327815	12/2021	Koponen et al.	N/A	N/A
11425126	12/2021	Horal et al.	N/A	N/A
11494518	12/2021	Curtis	N/A	G06F 21/629
11496517	12/2021	Hinrichs et al.	N/A	N/A
11509658	12/2021	Kulkarni	N/A	N/A
11604684	12/2022	Hinrichs et al.	N/A	N/A
11681568	12/2022	Hinrichs et al.	N/A	N/A
11741244	12/2022	Sandall et al.	N/A	N/A
11762712	12/2022	Koponen et al.	N/A	N/A
11847241	12/2022	Cahill et al.	N/A	N/A
11853463	12/2022	Hinrichs et al.	N/A	N/A
12118102	12/2023	Sandall et al.	N/A	N/A
2003/0115484	12/2002	Moriconi et al.	N/A	N/A
2003/0220925	12/2002	Lior	N/A	N/A
2004/0083367	12/2003	Garg et al.	N/A	N/A
2005/0081058	12/2004	Chang et al.	N/A	N/A
2005/0114674	12/2004	Carley	713/182	G06F 21/40
2006/0053290	12/2005	Randle	713/169	G06Q 20/027
2006/0059569	12/2005	Dasgupta et al.	N/A	N/A
2007/0006325	12/2006	Gargaro	N/A	N/A
2007/0156670	12/2006	Lim	N/A	N/A
2007/0226320	12/2006	Hager	709/219	G06F 16/137
2008/0022357	12/2007	Agarwal et al.	N/A	N/A
2008/0184336	12/2007	Sarukkai et al.	N/A	N/A
2009/0019533	12/2008	Hazlewood et al.	N/A	N/A
2009/0055749	12/2008	Chatterjee	715/738	G06F 3/04817
2009/0063665	12/2008	Bagepalli et al.	N/A	N/A
2009/0077618	12/2008	Pearce et al.	N/A	N/A
2009/0138960	12/2008	Felty et al.	N/A	N/A

2009/0281996	12/2008	Liu et al.	N/A	N/A
2010/0095373	12/2009	Levenshteyn et al.	N/A	N/A
2010/0333079	12/2009	Sverdlov et al.	N/A	N/A
2011/0113484	12/2010	Zeuthen	726/19	H04L 63/102
2012/0030354	12/2011	Razzaq et al.	N/A	N/A
2012/0066487	12/2011	Brown	713/150	G06F 9/5083
2012/0066756	12/2011	Vysogorets et al.	N/A	N/A
2012/0311672	12/2011	Connor et al.	N/A	N/A
2012/0331539	12/2011	Matsugashita	N/A	N/A
2013/0226970	12/2012	Weber et al.	N/A	N/A
2013/0227636	12/2012	Bettini et al.	N/A	N/A
2013/0283370	12/2012	Vipat et al.	N/A	N/A
2013/0305354	12/2012	King et al.	N/A	N/A
2014/0032691	12/2013	Barton et al.	N/A	N/A
2014/0032733	12/2013	Barton et al.	N/A	N/A
2014/0032759	12/2013	Barton et al.	N/A	N/A
2014/0033267	12/2013	Aciicmez	N/A	N/A
2014/0181186	12/2013	Stevens	709/203	H04L 67/322
2014/0237594	12/2013	Thakadu et al.	N/A	N/A
2014/0372986	12/2013	Levin	717/126	G06F 9/44589
2015/0089575	12/2014	Vepa et al.	N/A	N/A
2015/0213449	12/2014	Morrison et al.	N/A	N/A
2015/0244724	12/2014	Xu et al.	N/A	N/A
2016/0034900	12/2015	Nelsen et al.	N/A	N/A
2016/0057027	12/2015	Hinrichs et al.	N/A	N/A
2016/0057107	12/2015	Call et al.	N/A	N/A
2016/0103870	12/2015	Patiejunas et al.	N/A	N/A
2016/0188898	12/2015	Karinta et al.	N/A	N/A
2016/0205101	12/2015	Verma et al.	N/A	N/A
2016/0352695	12/2015	Kozolchyk et al.	N/A	N/A
2016/0373455	12/2015	Shokhrin	N/A	G06F 21/629
2016/0381032	12/2015	Hashmi et al.	N/A	N/A
2017/0024428	12/2016	Patiejunas et al.	N/A	N/A
2017/0075938	12/2016	Black et al.	N/A	N/A
2017/0111336	12/2016	Davis et al.	N/A	N/A
2017/0124166	12/2016	Thomas et al.	N/A	N/A
2017/0142068	12/2016	Devarajan et al.	N/A	N/A
2017/0161120	12/2016	Sasaki et al.	N/A	N/A
2017/0220370	12/2016	Klompje et al.	N/A	N/A
2017/0237729	12/2016	Uppalapati	N/A	N/A
2017/0279805	12/2016	Diaz-Cuellar et al.	N/A	N/A
2017/0302655	12/2016	Sondhi	N/A	H04L 63/10
2017/0331629	12/2016	Kozolchyk et al.	N/A	N/A
2017/0346807	12/2016	Blasi	N/A	N/A
2017/0364702	12/2016	Goldfarb et al.	N/A	N/A
2018/0062858	12/2017	Xu et al.	N/A	N/A
2018/0067790	12/2017	Chheda et al.	N/A	N/A
2018/0082053	12/2017	Brown et al.	N/A	N/A
2018/0109538	12/2017	Kumar et al.	N/A	N/A
2018/0295036	12/2017	Krishnamurthy et al.	N/A	N/A

2018/0309746	12/2017	Blasi	N/A	N/A
2019/0007418	12/2018	Cook et al.	N/A	N/A
2019/0007443	12/2018	Cook et al.	N/A	N/A
2019/0020665	12/2018	Surcouf et al.	N/A	N/A
2019/0080103	12/2018	Hadzic et al.	N/A	N/A
2019/0190959	12/2018	Yuan	N/A	N/A
2019/0230130	12/2018	Beckman et al.	N/A	N/A
2019/0245862	12/2018	Kruse et al.	N/A	N/A
2019/0273746	12/2018	Coffing	N/A	N/A
2019/0386973	12/2018	Patwardhan et al.	N/A	N/A
2020/0007580	12/2019	Liderman et al.	N/A	N/A
2021/0029029	12/2020	Mehmedagic et al.	N/A	N/A
2021/0240550	12/2020	Hinrichs et al.	N/A	N/A
2021/0248017	12/2020	Hinrichs et al.	N/A	N/A
2023/0094066	12/2022	Hecht	726/21	G06F 21/552
2024/0004728	12/2023	Hinrichs et al.	N/A	N/A

OTHER PUBLICATIONS

Non-Published commonly Owned U.S. Appl. No. 16/050,123, filed Jul. 31, 2018, 56 pages, Styra, Inc. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 16/050,124, filed Jul. 31, 2018, 56 pages, Styra, Inc. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 16/050,127, filed Jul. 31, 2018, 56 pages, Styra, Inc. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 16/050,130, filed Jul. 31, 2018, 56 pages, Styra, Inc. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 16/050,136, filed Jul. 31, 2018, 57 pages, Styra, Inc. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 16/050,139, filed Jul. 31, 2018, 56 pages, Styra, Inc. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 16/050,143, filed Jul. 31, 2018, 56 pages, Styra, Inc. cited by applicant

Author Unknown, "API Best Practices Managing the API Lifecycle: Design, Delivery, and Everything in Between," Dec. 2016, 37 pages, Apigee, retrieved from <https://pages.apigee.com/rs/351-WXY-166/images/API-Best-Practices-ebook-2016-12.pdf>. cited by applicant

Costa, Jeff, "Improve API Performance with Caching," API Gateway, May 31, 2018, 18 pages, Akamai Developer, retrieved from <https://developer.akamai.com/blog/2018/05/31/improve-api-performance-caching>. cited by applicant

Moffett, Jonathan D., et al., "Policy Hierarchies for Distributed Systems Management," IEEE Journal on Selected Areas in Communications, Dec. 1993, 11 pages, vol. 11, IEEE, USA. cited by applicant

Non-Published commonly Owned U.S. Appl. No. 18/120,810, filed Mar. 13, 2023, 63 pages, Styra, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/239,714, filed Aug. 29, 2023, 68 pages, Styra, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/369,471, filed Sep. 18, 2023, 74 pages, Styra, Inc. cited by applicant

Win, Thu Yein, et al., "Virtualization Security Combining Mandatory Access Control and Virtual Machine Introspection," 2014 IEEE/ACM 7th International Conference on Utility and Cloud

Computing, Dec. 8-11, 2014, 6 pages, IEEE, London, UK. cited by applicant
Preuveneers, Davy, et al., "Access Control with Delegated Authorization Policy Evaluation for Data-Driven Microservice Workflows," Future Internet, Sep. 30, 2017, 21 pages, vol. 9, Multidisciplinary Digital Publishing Institute. cited by applicant
Wei, Hao, et al., "Enhance OpenStack Access Control via Policy Enforcement Based on XACML," In Proceedings of the 16th International Conference on Enterprise Information Systems, Apr. 2014, 7 pages, vol. 1, SciTePress, Lisbon, Portugal. cited by applicant

Primary Examiner: Anya; Charles E

Attorney, Agent or Firm: ADELI LLP

Background/Summary

BACKGROUND

(1) Most applications today use access control rules to allow different users with different privileges, different access to the applications and their resources. Typically, the access controls are coded within an application's code base, which makes it very difficult to modify these controls statically, and impossible to modify them dynamically while the application is running. This problem is especially acute for micro-service applications, which are often small, well-defined applications that perform a small set of operations. Incorporating complex access controls in these applications is not practical as such controls would burden these applications with unnecessary functional complexity that not only runs counter to the focus of these micro-service applications, but would also be difficult to maintain and update. Accordingly, most micro-service applications employ no access controls or the most rudimentary access controls.

BRIEF SUMMARY

(2) Some embodiments of the invention provide a system for defining, distributing and enforcing policies for authorizing API (Application Programming Interface) calls to applications executing on one or more sets of associated machines (e.g., virtual machines, containers, computers, etc.) in one or more datacenters. This system has a set of one or more servers that acts as a logically centralized resource for defining and storing policies and parameters for evaluating these policies.

(3) The server set in some embodiments receives and stores definitions of several authorization policies for several API calls to one or more applications executing on a set of associated machines. In some embodiments, the sever set can receive the policies through a web-based user interface (e.g., a web browser) or some other interface (e.g., an API interface that supports receiving policies from other applications). The server set in some embodiments allows these policies to be custom defined in terms of different sets of parameters for the different sets of applications that are used in the different deployment settings in which the server set is employed (e.g., for different tenants in a multi-tenant datacenter or different departments or divisions of a multi-department or multi-division entity).

(4) Through one or more interfaces (e.g., web-based interfaces or other interfaces), the server set in some embodiments receives parameters for evaluating the policies. For one or more sources (e.g., LDAP directories, etc.) in the datacenter(s), the server set in some embodiments has one or more data-source adapters that collect parameters from these data sources and store these parameters in order to evaluate the authorization policies to assess whether API calls should be authorized. The server set also allows such parameters to be entered through a web-based user interface.

(5) For each set of related machines (e.g., for each tenant in a multi-tenant datacenter), the server set in some embodiments stores the defined authorization policies, and collected parameters needed

to resolve these policies, in a single hierarchical storage structure that can be accessed as a single runtime storage construct. The server set in some embodiments includes several modules for maintaining the policy and parameters in this runtime storage structure up to date. For instance, as mentioned above, the server set in some embodiments has one or more data-source adapters that receive and/or collect policy-resolving parameters for storage in the runtime hierarchical storage structure.

(6) In some embodiments, the hierarchical storage structure allows each particular policy and a set of collected parameters associated with the particular policy to be retrieved by providing a location identifier (e.g., a path) that identifies a location in the hierarchical storage structure that stores the particular policy. In some embodiments, this location also stores the set of collected parameters for the particular policy. In other embodiments, this location either stores the set of collected parameters for the particular policy, or stores one or more other location-identifiers that identify one or more other locations in the hierarchical storage structure at which the set collected parameters is previously stored in the structure. For each set of related machines (e.g., for each tenant in a multi-tenant datacenter, for each divisional of a multi-division company, for each department of a school, etc.), the hierarchical storage structure in some embodiments is a namespace that stores the policies and the parameters for resolving these policies for the set of related machines.

(7) In some embodiments, the single storage structure (e.g., the namespace) stores the policy for authorizing each API call as one or more lines of opcode (i.e., as one or more opcode expressions/statements). Each such policy is referred to as policy opcode in the discussion below. Also, policy operands in the discussion below refer to the parameters that are needed to resolve the policy opcodes for the API calls. The single storage structure in some embodiments stores such policy operands along with the policy opcodes. The policy opcodes in some embodiments are written in a declarative language. Also, in some embodiments, the policy operands are written in a language-independent, structured format, such as JSON (Javascript Object Notation) format or YAML (initially, “Yet Another Markup Language” and later, YAML Ain’t Markup Language). Typically, such formats (e.g., JSON, YAML, etc.) are human readable. Alternatively or conjunctively, in some embodiments, the policy operands are expressed in other formats, e.g., binary format, etc.

(8) The server set in some embodiments identifies updates to one or more parameter sets (e.g., to one or more operands) for one or more policies (e.g., policy opcodes), and stores these updates in the namespace. The server set in different embodiments employs different techniques for updating the policy operands. For instance, in some embodiments, the server set identifies an operand update by retrieving the update from a source of the operand (i.e., the parameter set) through a data pull operation and/or by receiving the update from the source through a data push operation.

(9) The server set in some embodiments updates the policy operands dynamically in real time. Accordingly, in some embodiments, the server set collects and updates parameters for resolving policies dynamically and in a deployment specific manner (i.e., in a manner that can be context specific for each set of associated machines for which the server set is deployed). The dynamic and deployment-specific way that the server set collects and updates policy-resolving parameters greatly enhances the configurability of the sever set in each deployment. This is because server set not only allows custom policies to be defined for each deployment, but also allows these policies to be resolved based on conditions and dynamic parameters that are specific to each deployment.

(10) In some embodiments, the server set is not only a logically-centralized resource for defining and storing policies and parameters, but is also a logically-centralized resource for distributing the defined policies and parameters to policy-enforcing local agents that execute near applications that are to process the API calls (e.g., to local agents that execute on the same host computers as the API-processing applications, on the same VM/container as the API-processing applications, etc.).

In some of the embodiments that stores the policies and parameters in a namespace, the server set distributes different portions of the namespace to different host computers based on the policies that are relevant to the API-authorization processing for the applications executing on the different host computers. The different portions that are distributed to the different host computers can be overlapping in some embodiments.

(11) For instance, to a first computer, the server set distributes a first portion of the namespace that comprises policies and parameter sets relevant for processing API calls to a set of applications executing on the first computer, while distributing a second portion of the namespace that comprises policies and parameter sets relevant for processing API calls to a set of applications executing on the second computer, with the first and second namespace portions differing in that at least one policy and at least one parameter set of the first portion is not in the second portion. These two namespaces can have no policies or parameters in common, or they can have one or more policies and parameters in common.

(12) In some embodiments, the server set uses a push model to distribute the policies/parameters (e.g., the relevant namespace portions) to each local agent. In other embodiments, the server set uses a pull model to distribute the policies/parameters. For instance, in these embodiments, once a local agent initializes, it queries the server set for policies and parameters relevant to it. In these embodiments, the local agent also periodically polls the server set to determine whether policies or parameters that the local agent previously received have been updated on the server set, and thereby need to be copied on the local agent's portion of the namespace. When such updates exist, the local agent retrieves them from the server set or receives them as part of its querying poll.

(13) The local agent for an application receives requests to determine whether API calls received by the application are authorized. The local agent in some embodiments receives each request to authorize an API call to an application from an authorization module of that application. In some embodiments, the authorization module of an application sends its API-authorization requests to the local agent through an RPC (Remote Procedure Call) message that is forwarded to the local agent by a network communication stack (e.g., the network stack executing on the application's host computer, the application's VM, or the application's container).

(14) In other embodiments, the authorization module sends its API-authorization requests to the local agent through other mechanisms (such as other IPC (inter-process communication) messages) that initiate above the network stack. Also, in some embodiments, the API-authorization request does not come to the local agent from the application to which the API call is directed. For instance, in some embodiments, a first application intercepts an API call intended for a second application, and sends a request (e.g., through RPC, IPC, etc.) to authorize this API call to the local agent. When the local agent authorizes this request, the first application then passes the API call to the second application, while discarding this call when the local agent rejects this call.

(15) The received API-authorization request contains a first set of one or more parameters associated with the API call. The first parameter set includes at least one parameter that identifies the API call, its type and/or other metadata about the API call. In some embodiments, at least one parameter in the first parameter set is received by the application in the API call. For instance, in some embodiments, the application receives the API call as part of the payload of one or more data messages (e.g., data packets) that the application receives, and at least one parameter in the first parameter set is a parameter that is part of the payload of the received data message(s). One or more modules (i.e., processes) are typically used to extract this payload for processing. For example, a layer 4 security process (e.g., a transport layer security (TLS) process) might have to decrypt one or more packet payloads in order to obtain the API call and its parameters in plaintext format (as opposed to the cipher text format in which they are received).

(16) After receiving the API-authorization request from a particular application, the local agent uses one or more parameters in the first parameter set to retrieve policy opcode related to the API call from the namespace that it received from the server set. For a typical local agent, the agent's

namespace is a subset of the global namespace that the server set maintains for a set of related machines, because the local agent's namespace in some embodiments only contains the policies that are relevant for the API calls that the agent will encounter, as mentioned above.

(17) To evaluate this policy opcode, the agent might also retrieve policy operands, which serves as a second set of parameters associated with the API call. The local agent then uses the first parameter set and/or second parameter set to evaluate the retrieved API-authorization policy opcode in order to determine whether the API call should be approved or rejected. Specifically, the policy opcode includes one or more rules that specify one or more conditions for approving or rejecting the API calls. The agent resolves these conditions based on the first parameter set and/or second parameter set.

(18) For instance, in some embodiments, the agent determines that one or more parameters in the received parameter set and/or retrieved parameter set match a condition specified by the policy opcode (in terms of one or more parameters) for rejecting an API call. In some embodiments, the conditions can specify reasons for allowing an API call. Some API-authorization policies are resolved only by reference to the received parameter sets (i.e., only specify conditions that just need the received parameter sets to be resolved). Other API-authorization policies are resolved only by reference to the retrieved parameter sets (i.e., only specify conditions that just need the retrieved parameter sets to be resolved). Still other API-authorization policies are resolved by reference to both the received and retrieved parameter sets (i.e., specify conditions that need both the received and retrieved parameter sets to be resolved).

(19) Based on this evaluation, the local agent then sends the particular application (e.g., the API authorization module of the particular application) a message (e.g., an RPC reply message) to specify whether the API call should be approved or rejected. When the API call is approved, the particular application then performs the operation associated with the API call, and if necessary, returns an output to the source of the API call (i.e., to the module that sent the API call) to the particular application. On the other hand, when the API call is rejected, the particular application sends the message to the source of the API call to reject the call, or simply discards this call without providing a response.

(20) In some embodiments, the local agents operate in the same failure domain as the machines on which the applications execute. For instance, when an application executes on a virtual machine (VM) or container, the local agent in some embodiments executes on the same VM or container as the application. In other embodiments, the local agent does not execute on the same VM or container as its associated application, but executes on the same host computer as the application's VM or container.

(21) In some embodiments, each local agent is associated with just one application that executes on the same computer, same VM or same container as the local agent. In other embodiments, one local agent in some embodiments can be shared by multiple applications that execute on the same VM, container, or host computer as the agent. However, in a multi-tenant datacenter, some embodiments do not allow local agents to be shared by applications executing on different tenant computers, VMs or containers. Other embodiments, on the other hand, allow one local agent to be used by different applications of different tenants.

(22) Having local agents authorize API calls is very useful for micro-service applications as it provides a mechanism to implement complex API-authorization policies without adding to the micro-service applications the functional complexity needed to process rich API-authorization policies. This approach is also useful as it extracts the API-authorization processing from the application's code base, and therefore makes this processing easier to change dynamically over time. It also allows API-authorization processing for a large number of applications in a datacenter to be controlled in a logically centralized way through the server set.

(23) As mentioned above, the server set in some embodiments distributes policies and parameters to local agents so that they can evaluate API-authorization requests locally. Conjunctively or

alternatively, the server set in some embodiments also processes API-authorization requests from applications executing on a set of associated machines. For instance, in some embodiments, it might not be feasible or desirable to perform the API-authorization on the same computer or failure domain as the application, e.g., the operating system restricts the type of agents installed, or the API-processing application executes on devices with limited computational resources (e.g., IoT (Internet of Things) devices). Also, in some embodiments, the response time for the API-authorization decision is not as crucial as the requested operations can be gradually rejected, or a limited access to the application can be provided until the API calls are authorized.

(24) Accordingly, in some embodiments, an API-authorization module of an application executing on a computer sends an API-authorization request (e.g., as an RPC message) to a local agent on the same computer, and this local agent forwards this request (e.g., as another RPC message) to the server set. In other embodiments, the API-authorization module of the application sends the API-authorization request (e.g., as an RPC message) directly to the server set.

(25) The received API-authorization request contains a first parameter set associated with the API call. As before, the first parameter set includes at least one parameter that identifies the API call, its type and/or other metadata about the API call, and can include one parameter that was received as part of the API call. After receiving the API-authorization request, the server set then uses one or more parameters in the first parameter set to retrieve a policy opcode related to the API call from the namespace for the set of associated machines to which the application belongs. To evaluate this policy opcode, the server set might also retrieve policy operands from the namespace, which serves as a second set of parameters associated with the API call.

(26) The server set then uses the first parameter set and/or second parameter set to evaluate the retrieved API-authorization policy opcode in order to determine whether the API call should be approved or rejected. For example, in some embodiments, the server set resolves one or more conditions specified in the retrieved policy opcode based on one or more parameters in the first parameter set and/or second parameter set (e.g., determines that one or more parameters in the received parameter set and/or retrieved parameter set matches a parameter-related condition specified by the policy opcode for rejecting an API call). Based on this evaluation, the server set then sends the local agent or the application's authorization module a reply (e.g., a RPC reply message) to specify whether the API call should be approved or rejected. When this response is sent to the local agent, the agent then relays this message (e.g., as an RPC reply message) to the application's API-authorization module.

(27) In some embodiments, the API-authorization modules (on the host computers or on the server set) are used to authorize API calls for both application domain logic (also called core logic or business logic) and application infrastructure logic. In other embodiments, the API-authorization modules (on the host computers or on the server set) are used to authorize API calls for application infrastructure logic but not application domain logic.

(28) Applications typically include application domain logic and application infrastructure logic. Application domain logic captures the domain rules for creating, storing, and manipulating data relating to the core functions of the application. For instance, in banking, domain logic captures the logic to calculate the monthly interests for bank accounts, prepare and calculate a mortgage offer (involving complex formulas using credit scores and other information available), or simply move funds between two accounts. In banking, domain logic would also be responsible for preparing the web page contents as customers access their bank accounts over the web.

(29) Application infrastructure logic provides support functions that allow the application domain logic to run, and to run reliably and scalably. Examples of such support functions include integration functions with other applications (various protocol implementations), configuration management (logic might have configuration options), monitor and alert functions (to detect when the domain logic is down), schedule and thread management facilities (to run several logic instances in parallel for improved scalability), etc. While these functions are critical, they are not

typically a concern of the application domain logic. Hence, decoupling API authorization for infrastructure logic simplifies the application code and allows the application's code to focus on matters relevant to the application's core functions.

(30) The preceding Summary is intended to serve as a brief introduction to some embodiments of the invention. It is not meant to be an introduction or overview of all inventive subject matter disclosed in this document. The Detailed Description that follows and the Drawings that are referred to in the Detailed Description will further describe the embodiments described in the Summary as well as other embodiments. Accordingly, to understand all the embodiments described by this document, a full review of the Summary, Detailed Description, the Drawings, and the Claims is needed. Moreover, the claimed subject matters are not to be limited by the illustrative details in the Summary, Detailed Description, and the Drawings.

Description

BRIEF DESCRIPTION OF FIGURES

- (1) The novel features of the invention are set forth in the appended claims. However, for purposes of explanation, several embodiments of the invention are set forth in the following figures.
- (2) FIG. 1 illustrates a policy enforcement system of some embodiments of the invention.
- (3) FIG. 2 illustrates the host-side software architecture of the API-authorization system of some embodiments.
- (4) FIG. 3 illustrates a process that an application performs when it receives an API call.
- (5) FIG. 4 illustrates a more-detailed view of a local agent of some embodiments.
- (6) FIG. 5 illustrates a process that a local agent performs when it receives a request to authorize an API call from an API-authorizing agent of an application in some embodiments.
- (7) FIGS. 6-8 present three examples that illustrate resolving the same API authorization request based on deployment specific policies and operands with the values of the operands changing dynamically based on network conditions.
- (8) FIG. 9 illustrates an example of the policy opcode and operand that can be used to implement API-authorization controls.
- (9) FIG. 10 illustrates the software architecture of the server set of some embodiments of the invention.
- (10) FIG. 11 conceptually illustrates an electronic system with which some embodiments of the invention are implemented.

DETAILED DESCRIPTION

- (11) In the following detailed description of the invention, numerous details, examples, and embodiments of the invention are set forth and described. However, it will be clear and apparent to one skilled in the art that the invention is not limited to the embodiments set forth and that the invention may be practiced without some of the specific details and examples discussed.
- (12) Some embodiments of the invention provide a system for defining, distributing and enforcing policies for authorizing API (Application Programming Interface) calls to applications executing on one or more sets of associated machines (e.g., virtual machines, containers, computers, etc.) in one or more datacenters. This system has a set of one or more servers that acts as a logically centralized resource for defining, storing, and distributing policies and parameters for evaluating these policies.
- (13) The server set in some embodiments receives and stores definitions of several authorization policies for several API calls to one or more applications executing on a set of associated machines. In some embodiments, the sever set can receive the policies through a web-based user interface (e.g., a web browser) or some other interface (e.g., an API interface that supports receiving policies from other applications). The server set in some embodiments allows these policies to be custom defined in terms of different sets of parameters for the different sets of applications that are used in

the different deployment settings in which the server set is employed (e.g., for different tenants in a multi-tenant datacenter or different departments or divisions of a multi-department or multi-division entity).

(14) Through one or more interfaces (e.g., web-based interfaces or other interfaces), the server set in some embodiments collects and stores parameters for evaluating the authorization policies to assess whether API calls should be authorized. For one or more sources (e.g., LDAP directories, etc.) in the datacenter(s), the server set in some embodiments has one or more data-source adapters that collect parameters from these data sources and store these parameters in order to evaluate the authorization policies to assess whether API calls should be authorized. The server set also allows such parameters to be entered through a web-based user interface.

(15) The collected parameters in some embodiments are specific to each set of related machines (e.g., for each tenant in a multi-tenant datacenter or each department or division of a multi-department or multi-division entity) for which the server set defines and distributes policies. The server set in some embodiments updates these parameters dynamically in real time. Accordingly, in some embodiments, the server set uses, collects and updates parameters for resolving policies dynamically and in a deployment specific manner (i.e., in a manner that can be context specific for each set of associated machines for which the server set is deployed). The dynamic and deployment-specific way that the server set collects and updates policy-resolving parameters greatly enhances the configurability of the server set in each deployment. This is because the server set not only allows custom policies to be defined for each deployment based on the deployment-setting specific parameters, but also allows these policies to be resolved based on dynamically changeable parameters for each specific deployment.

(16) For each set of related machines, the server set in some embodiments stores the defined API-authorization policies, and collected parameters needed to resolve these policies, in a single hierarchical storage structure (such as a namespace) that can be accessed as a single runtime storage construct. The server set in some embodiments includes several modules for maintaining the policy and parameters in this runtime storage structure up to date. For instance, as mentioned above, the server set in some embodiments has one or more data-source adapters that receive and/or collect policy-resolving parameters for storage in the runtime hierarchical storage structure.

(17) The server set in some embodiments identifies updates to parameters for evaluating the policies and stores these updates in this storage structure along with the authorization policies. In some embodiments, the runtime storage structure (e.g., the namespace) stores the policy for authorizing each API call as one or more lines of opcode (as one or more opcode expressions/statements). Each such policy is referred to as policy opcode in the discussion below. Also, policy operands in the discussion below refer to the parameters that are needed to resolve the policy opcodes for the API calls. The single storage structure stores such policy operands along with the policy opcodes. Policy operands can also accompany the API calls, or the API-authorization requests that are generated in response to such API calls.

(18) In some embodiments, the server set distributes the defined policies and parameters to policy-enforcing local agents that execute near the applications that process the API calls. Different embodiments place the local agents in different positions with respect to their associated applications, as further described below. In some embodiments, the server set distributes different sets of policies and parameters to different local agents based on the policies that are relevant to the API-authorization processing for the applications associated with those agents.

(19) From an associated application, a local agent receives API-authorization requests to determine whether API calls received by the application are authorized. In response to such a request, the local agent uses one or more parameters associated with the API call to identify a policy stored in its local policy storage to evaluate whether the API call should be authorized. To evaluate this policy, the agent might also retrieve one or more parameters from the local policy storage.

(20) The local agent then uses the received or retrieved parameters to evaluate the retrieved API-

authorization policy in order to determine whether the API call should be approved or rejected. Based on this evaluation, the local agent then sends a reply to the application that sent the API-authorization request, to specify whether the API call should be approved or rejected. In some embodiments, the applications send the API-authorization requests as RPC (Remote Procedure Call) messages, and the local agents send their replies to these requests as RPC reply messages.

(21) In other embodiments, the applications send their API-authorization requests to the local agents through other mechanisms (such as other IPC (inter-process communication) messages) that initiate above the network stack. Also, in some embodiments, an API-authorization request does not come to the local agent from the application to which the API call is directed. For instance, in some embodiments, a first application intercepts an API call intended for a second application, and sends a request (e.g., through RPC, IPC, etc.) to authorize this API call to the local agent. When the local agent authorizes this request, the first application then passes the API call to the second application, while discarding this call when the local agent rejects this call. In some embodiments, the first application not only intercepts API calls to obtain authorization for these calls, but also intercepts these calls to first have them authenticated by a separate authentication service.

(22) In some embodiments, the local agents are used to authorize API calls for both application domain logic (also called core logic or business logic) and application infrastructure logic. In other embodiments, the local agents are used to authorize API calls for application infrastructure logic but not application domain logic. Applications typically include application domain logic and application infrastructure logic. Application domain logic captures the domain rules for creating, storing, and manipulating data relating to the core functions of the application. For instance, in banking, domain logic captures the logic to calculate the monthly interests for bank accounts, prepare and calculate a mortgage offer (involving complex formulas using credit scores and other information available), or simply move funds between two accounts. In banking, domain logic would also be responsible for preparing the web page contents as customers access their bank accounts over the web.

(23) Application infrastructure logic provides support functions that allow the application domain logic to run, and to run reliably and scalably. Examples of such support functions include integration functions with other applications (various protocol implementations), configuration management (logic might have configuration options), monitor and alert functions (to detect when the domain logic is down), schedule and thread management facilities (to run several logic instances in parallel for improved scalability), etc. While these functions are critical, they are not typically a concern of the application domain logic. Hence, decoupling API authorization for infrastructure logic simplifies the application code and allows the application's code to focus on the matters relevant to the application's core functions.

(24) FIG. 1 illustrates a policy enforcement system **100** of some embodiments of the invention. As shown, this system includes a server set **105**, several host computers **110**, a first set of applications **115** executing on the host computers **110**, several local agents **120** executing on the host computers, and a second set of applications **125** executing on the same host computers **110** or other devices or computers **130**. The various computers and devices of the system **100** communicatively connect to each other through a network **180**, which can include one or more local area networks, wide area networks, wireless networks, a network of networks (e.g., Internet), etc.

(25) The host computers **110** in some embodiments operate in one or more datacenters. In some embodiments, the server set **105** includes one or more servers that operate in the same datacenter as the host computers **110**, while in other embodiments the server set operates in a different datacenter than one or more of the host computers. The server set **105** acts as a logically centralized resource for defining, storing, distributing and enforcing policies for authorizing API calls to the first set of applications **115** that execute sets of associated machines (e.g., VMs, containers, etc.) that execute on the host computers **110**. In some embodiments, each set of associated machines is associated with one entity (e.g., belongs to a tenant in a multi-tenant datacenter or to one division of a

corporation in a corporate datacenter). Accordingly, the server set can be used to define, store, distribute and enforce different sets of policies for authorizing API-calls for the set of applications **115** that belong to different entities (e.g., different tenants) in one or more datacenters.

(26) In some embodiments, some of these applications execute on virtual machines (VMs) or containers on the host computers, while other applications execute over the operating systems of the host computers. Also, in some embodiments, some of the API calls to a set of associated applications **115** come from other applications in the same set, while other API calls to the application set come from other applications (e.g., applications **125**). The second set of applications **125** execute on host computers or devices inside or outside of the server set's datacenter(s) for which the server set does not enforce policies.

(27) To define and store policies, the server set **105** in some embodiments receives and stores definitions of several authorization policies for a number of API calls to a set of associated applications. In some embodiments, the sever set can receive the policies through a web-based user interface (e.g., a web browser) or some other interface (e.g., an API interface that supports receiving policies from other applications). The server set in some embodiments allows these policies to be custom defined in terms of different sets of parameters for the different sets of applications that are used in the different deployment settings in which the server set is employed (e.g., for different tenants in a multi-tenant datacenter or different departments or divisions of a multi-department or multi-division entity).

(28) Through one or more interfaces (e.g., web-based interfaces or other interfaces), the server set **105** in some embodiments receives parameters for evaluating the policies. For one or more policy operand data sources **132** (e.g., from tenant LDAP directories, etc.) in the datacenter(s), the server set in some embodiments has one or more adapters that collect parameters from these data sources and store these parameters in order to evaluate the authorization policies to assess whether API calls should be authorized. The server set also allows such parameters to be entered through a web-based user interface.

(29) For each set of related applications or machines (e.g., for each tenant in a multi-tenant datacenter), the server set in some embodiments stores the defined authorization policies and collected parameters in a single hierarchical storage structure **135**. In some embodiments, the hierarchical storage structure allows each particular policy and a set of collected parameters associated with the particular policy to be retrieved by providing a location identifier (e.g., a path) that identifies a location in the hierarchical storage structure that stores the particular policy. In some embodiments, this location also stores the set of collected parameters for the particular policy. In other embodiments, this location either stores the set of collected parameters for the particular policy, or stores one or more location-identifiers for one or more locations in the hierarchical storage structure at which the set collected parameters is previously stored in the structure. The hierarchical storage structure **135** in some embodiments is a namespace that stores the policies and the parameters for resolving these policies for the set of related machines.

(30) In some embodiments, each authorization policy is defined by policy opcode in the namespace, while the parameter set for this authorization policy is defined as policy operands that are associated with the policy's opcode and that are stored in the namespace. The policy opcodes in some embodiments are written in a declarative language. Also, in some embodiments, the policy operands are written in a language-independent, structured format, such as JSON (Javascript Object Notation) format or YAML (initially, "Yet Another Markup Language" and later, YAML Ain't Markup Language). Typically, such formats (e.g., JSON, YAML, etc.) are human readable. Alternatively or conjunctively, in some embodiments, the policy operands are expressed in other formats, e.g., binary format, etc.

(31) The server set **105** in some embodiments identifies updates to one or more parameter sets (e.g., to one or more operands) for one or more one policies (e.g., policy opcodes), and stores these updates in the namespace. The server set in different embodiments employs different techniques for

updating the policy operands. For instance, in some embodiments, the data source adapters of the server set identify operand updates by retrieving the updates from the data source **132** of the operands through data pull operations and/or by receiving the updates from these data source **132** through data push operations.

(32) In some embodiments, the server set distributes the defined policies and parameters to local agents **120** that enforce these API-authorization policies for the applications **115**, and these local agents **120** store the distributed policies and parameters in their local policy/parameter store **150**. In some of the embodiments that store the policies and parameters in a namespace, the server set distributes different portions of the namespace to different host computers based on the policies that are relevant to the API-authorization processing for the first set of applications **115** executing on the different host computers. The different portions that are distributed to the different host computers can be overlapping in some embodiments.

(33) For instance, to a first computer **110a**, the server set **105** distributes a first portion of the namespace that comprises policies and parameter sets relevant for processing API calls to a set of applications executing on the first computer, while distributing a second portion of the namespace that comprises policies and parameter sets relevant for processing API calls to a set of applications executing on the second computer **110b**, with the first and second namespace portions differing in that at least one policy and at least one parameter set of the first portion is not in the second portion. These two namespaces can either have no policies or parameters in common, or have one or more policies and parameters in common. Each local agent **120** stores its local portion of the namespace that it receives from the server set **105** in its local policy/parameter store **150**.

(34) In some embodiments, the server set **105** uses a push model to distribute the policy/parameters (e.g., the relevant namespace portions) to each local agent **120**. In other embodiments, the server set uses a pull model to distribute the policy/parameters. For instance, in these embodiments, once a local agent initializes, it queries the server set for policies and parameters relevant to it. In these embodiments, the local agent also periodically polls the server set to determine whether policies or parameters that the local agent previously received have been updated on the server set, and thereby need to be copied on the local agent's portion of the namespace. When such updates exist, the local agent retrieves them from the server set or receives them as part of its querying poll.

(35) The local agent **120** executing on a host computer **110** receives requests to determine whether API calls received by one or more applications **115** executing on the host computer are authorized. The local agent in some embodiments receives each request to authorize an API call to an application from an authorization module of that application. In some embodiments, the authorization module of an application sends its API-authorization requests to the local agent through an RPC (Remote Procedure Call) message that is forwarded to the local agent by a network communication stack (e.g., the network stack executing on the application's host computer, the application's VM, or the application's container). In other embodiments, the authorization module sends its API-authorization requests to the local agent through other mechanisms (such as other IPC (inter-process communication) messages) that again, initiate above the network stack.

(36) In response to such a request, the local agent **120** uses one or more parameters associated with the API call to identify a policy (e.g., policy opcode) stored in its local policy storage (e.g., stored in its local namespace) to evaluate in order to determine whether the API call should be authorized. The received API-authorization request contains one or more parameters associated with the API call (e.g., one or more parameters that identify the API call, its type and/or other metadata about the API call). In some embodiments, at least one parameter in the request is received by the application in the API call.

(37) For instance, in some embodiments, the application receives the API call as part of the payload of one or more data messages (e.g., data packets) that the application receives, and at least one parameter in the first parameter set is a parameter that is part of the payload of the received data message(s). One or more modules (i.e., processes) are typically used to extract this payload for

processing. For example, a layer **4** security process (e.g., a transport layer security (TLS) process) might have to decrypt one or more packet payloads in order to obtain the API call and its parameters in plaintext format (as opposed to the cipher text format in which they are received). (38) To evaluate the identified policy for the API call, the agent might also retrieve one or more parameters (e.g., operands) from the local policy storage (e.g., from its copy of the namespace). The local agent then uses the received or retrieved parameters to evaluate the retrieved API-authorization policy in order to determine whether the API call should be approved or rejected. For instance, in some embodiments, the retrieved policy opcode includes one or more rules that specify one or more conditions for approving or rejecting the API calls. The agent resolves these conditions based on the received or retrieved parameter(s) associated with the API call (e.g., the agent determines that one or more received or retrieved parameters match a condition specified by the policy opcode for rejecting an API call).

(39) Based on this evaluation, the local agent then sends a reply to the application that sent the API-authorization request, to specify whether the API call should be approved or rejected. In some embodiments, the local agent sends its reply to the API-authorization request as an RPC reply message. When the API call is approved, the particular application then performs the operation associated with the API call, and if necessary, returns an output to the source of the API call (i.e., to the module that sent the API call) to the particular application. On the other hand, when the API call is rejected, the particular application sends the message to the source of the API call to reject the call, or simply discards this call without providing a response to the source of the API call.

(40) In some embodiments, the local agents **120** operate in the same failure domain as the machines on which the applications **115** execute. For instance, when an application executes on a virtual machine (VM) or container, the local agent in some embodiments executes on the same VM or container as the application. In other embodiments, the local agent does not execute on the same VM or container as its associated application, but executes on the same host computer as the application's VM or container.

(41) In some embodiments, each local agent **120** is associated with just one application **115** that executes on the same computer, same VM or same container as the local agent. In other embodiments, one local agent **120** in some embodiments can be shared by multiple applications **115** that execute on the same VM, container, or host computer as the agent. However, in a multi-tenant datacenter, some embodiments do not allow local agents to be shared by applications executing on different tenant computers, VMs or containers. Other embodiments, on the other hand, allow one local agent to be used by different applications of different tenants.

(42) Some embodiments have the local agents **120** operate within the same failure domain as the applications (e.g., micro service applications) in order to ensure that the policy functionality remains highly available and responsive at all times. Having local agents authorize API calls is very useful for micro-service applications as it provides a mechanism to implement complex API-authorization policies without adding to the micro-service applications the functional complexity needed to process rich API-authorization policies. This approach is also useful as it extracts the API-authorization processing from the application's code base, and therefore makes this processing easier to change dynamically over time. It also allows API-authorization processing for a large number of applications in a datacenter to be controlled in a logically centralized way through the server set.

(43) As mentioned above, the server set **105** in some embodiments distributes policies and parameters to local agents **120** so that they can evaluate API-authorization requests locally. Conjunctively or alternatively, the server set **105** in some embodiments also processes API-authorization requests from some or all of the applications **115**. For instance, in some embodiments, it might not be feasible or desirable to perform the API-authorization on the same computer or failure domain as the application, e.g., the operating system restricts the type of agents installed, or the API-processing application executes on devices with limited computational

resources (e.g., IoT (Internet of Things) devices). Also, in some embodiments, the response time for the API-authorization decision is not as crucial as the requested operations can be gradually rejected, or a limited access to the application can be provided until the API calls are authorized. (44) Accordingly, in some embodiments, an API-authorization module of an application executing on a computer sends an API-authorization request (e.g., as an RPC message) to a local agent **120** on the same computer, and this local agent **120** forwards this request (e.g., as another RPC message) to the server set **105**. In other embodiments, the API-authorization module of the application sends the API-authorization request (e.g., as an RPC message) directly to the server set **105**.

(45) The received API-authorization request contains a first parameter set associated with the API call. As before, the first parameter set includes at least one parameter that identifies the API call, its type and/or other metadata about the API call, and can include one parameter that was received as part of the API call. After receiving the API-authorization request, the server set **105** then uses one or more parameters in the first parameter set to retrieve a policy (e.g., policy opcode) related to the API call from its policy/data storage **135** (e.g., from its namespace for the set of associated machines to which the application belongs). To evaluate this policy, the server set might also retrieve one or more parameters from its policy storage **135** that are relevant for assessing the retrieved policy (e.g., retrieve policy operands from its namespace for the set of associated machines to which the application belongs). These retrieved set of parameters serve as a second set of parameters associated with the API call.

(46) The server set **105** then uses the received or retrieved parameters to evaluate the retrieved API-authorization policy in order to determine whether the API call should be approved or rejected. For example, in some embodiments, the server set resolves one or more conditions specified in the retrieved policy opcode based on one or more received or retrieved parameters (e.g., determines that one or more received or retrieved parameters match a condition specified by the policy opcode for rejecting an API call). Based on this evaluation, the server set **105** then sends a reply (e.g., a RPC reply message) to the local agent **120** or the application's authorization module to specify whether the API call should be approved or rejected. When this response is sent to the local agent **120**, this agent then relays this message (e.g., as an RPC reply message) to the application's API-authorization module.

(47) As mentioned above, the local policy and data storage **150** and the server policy and data storage **135** of some embodiments are hierarchical namespaces that store both policy opcode and policy operands that are needed for resolving the policy opcodes (i.e., both the code and the data on which the code operates). In some embodiments, the namespace is a document database in that it has no tables and instead is laid out as a single hierarchical document. This document forms one global namespace that holds all the data and policies. For instance, the following namespace structure is a hierarchical layout of a policy-and-document database that stores three JSON data documents and two policies: /a/b/doc_1.json /a/b/doc_2.json /c/d/doc_3.json /d/policy_A /e/f/policy_B

(48) In this example, there are two JSON documents at namespace location '/a/b' and a third JSON document at '/c/d', while the two policies are saved at '/d' and '/e/f'.

(49) As a namespace merges the data and policies into one global hierarchical layout, database clients can access a nested portion of the namespace simply by requesting data by its fully qualified name. Moreover, to access a portion of policy operand document **1**, the client can simply request the data with name: /a/b/doc_1/some_portion.

The same applies to policies: the namespace evaluates the policies on-the-fly when client accesses the namespace. That is, client can simply access the results of a "policy evaluation" by referring the policy and its relevant portion. For instance, a client can request

/e/f/policy_B/some_computed_result,

which would make title database take the source of code of the policy B, compute the policy using

any data necessary that the policy may need underneath and then pass the requested portion of the policy results back to the client.

(50) FIG. 2 illustrates the host-side software architecture of the API-authorization system **100** of some embodiments. Specifically, it illustrates a host computer **200** on which multiple machines **205** (e.g., multiple VMs or containers) execute. On each machine **205**, multiple applications **210** can execute to process API calls from other modules on the host computer **200** or on other host computers or devices in the same datacenter or outside of the datacenter.

(51) In addition to these applications, each machine **205** also has a network stack **215** and a local API-authorizing agent **220**. The network stack **215** is a software communication stack (e.g., TCP/IP stack) that implements a network communication protocol suite to allow the applications and processes of the machine **205** to communicate with each other, and with applications and processes executing outside of this machine (e.g., on the same host computer **200** or on other computers or devices). When the machine **205** is a container, its network stack is the network stack of the VM or operating system on which the container executes. Also, the network stack of a machine often connects with another module (e.g., a software switch) on the host computer, which, in turn, interfaces with a physical network interface card (NIC) of the host computer.

(52) One set of communications that the network stack **215** facilitates is the communication between an application's API-authorizing agent **230** on an application and the local API-authorizing agent **220**. As further described above and below, the network stack passes RPC data messages between the application agent **230** and the local agent **220**. The local agent **220** enforces API-authorization policies for the applications that execute on the machine **205**. Specifically, from the application API-authorizing agents **230** of the machine **205**, the local API-authorizing agent **220** receives API-authorization requests for different APIs that the applications **210** receive.

(53) For each such request, the local agent **220** retrieves the API-authorization policy associated with the request (i.e., associated with API call to which the request pertains) from its local policy and parameter storage **225**. The local agent **220** retrieves this policy by using one or more parameters that are associated with the API call and that were received with the API call or the API-authorization request. From the policy/parameter storage, the local agent **220** can also retrieve one or more parameters for evaluating this policy.

(54) The local storage **225** stores policies and parameters that the server set **105** has previously provided to the local agent **220**. In some embodiments, this storage also stores parameters that the local agent previously receives from a local parameter integrator **255**. This integrator allows an entity in the datacenter (e.g., a tenant or a department) to separately provide some or all of the parameters for evaluating the API-authorization policies. For instance, this integrator allows the entity to export highly confidential parameters into an API-authorization policy enforcement node without providing these parameters to the server set **105**.

(55) The local agent **220** then evaluates the retrieved policy based on parameter(s) received with the API-authorization request, or retrieved from the local storage **225**. Based on its evaluation, the local agent then determines whether the API call should be approved or rejected. The local agent then sends the application API-authorizing agent **230** that sent the API-authorization request, a reply to specify whether the API call should be approved or rejected. In some embodiments, the API-authorizing agent **230** sends the API-authorization requests as RPC (Remote Procedure Call) messages, and the local agents send their replies to these requests as RPC reply messages.

(56) In the example illustrated in FIG. 2, each machine **205** has its own local API-authorization agent **220**. Each such local agent **220** on a machine **205** can process API-authorization requests from one or more applications **210** executing on the machine. Having the local agent **220** operate on the same machine as its associated application **210** is beneficial as in this way, both these components operate in the same failure domain, which improves agent's availability and its speed of processing API-authorization requests. However, in other embodiments, the applications on multiple machines on a host computer share the same local agent. For instance, in some

embodiments, the local agent operates as a service VM or container on a host computer, and applications on multiple VMs and/or containers use this agent to have their API calls authorized. Instead of having the local agent operate as a service VM or container in the user space, other embodiments have the local agent operate as another type of process on the host computer that is accessed by multiple applications on multiple VMs and/or containers on a host computer.

(57) As shown in FIG. 2, each application has an API handler **226**, an API-authentication agent **228**, and one or more API processing modules **224**, in addition to the API-authorizing agent **230**. An API processing module **224** is a module (e.g., an object) that performs an operation based on API call to it. The API authentication agent **228** is the module that authenticates the identity of the API-call's source (e.g., validates the access credentials associated with the API call). When the API call is authenticated, the API authorization agent **230** determines whether the API call is authorized (i.e., whether the operation associated with the API call should be performed). The API handler **226** is the module that initiates the authentication operation and authorization operation for each API call. In some embodiments, the API handler is implemented by authentication and/or authorization code of the API processing modules (e.g., this handler is implemented in some embodiments by authentication and authorization code within each API processing object of an application).

(58) To further explain the operations of the components of the application **210**, FIG. 3 illustrates a process that an application **210** performs when it receives (at **305**) an API call. In some embodiments, the application **210** performs this process only for an API call that relates to an application infrastructure operation. In other words, an application **210** does not call the local agent **220** in some embodiments for API calls that relate to an application's core logic operations. In other embodiments, the local agent **220** handles API-authorization requests for both core logic API calls and infrastructure logic API calls.

(59) When an API call is received, it has to be authenticated in some embodiments to validate the identity of the source of the API call. The API handler **226** uses the API authentication agent **228** to authenticate the identity of the source of the API call, e.g., to authenticate the credential set (e.g., certificate, username and/or password) accompanying the API call. Accordingly, as shown, the API handler **226** initially determines (at **310**) whether the received API call is part of a communication session that was previously authenticated by the authentication agent **228** (e.g., part of a communication session that the authentication agent **228** previously stored a cookie). If so, the process transitions to **325**, which will be further described below. Otherwise, when the API call is not part of a communication session that was previously authenticated, the API handler **226** directs (at **315**) the authentication agent **228** to authenticate the API call. In some embodiments, each API call has to be independently authenticated irrespective of whether any prior API call in the same communication session has been authenticated.

(60) In some embodiments, the API authentication agent **228** authenticates (at **315**) the identity of the source of the API call by authenticating the credential set that is part of the API call. The application receives an API call as part of one or more data messages, each of which has a header and a payload. The payload of each such message contains all or part of the API call and one or more associated parameters, e.g., parameters that provide the credentials (e.g., username, password, user identifier, group identifier, certificate, etc.) of the source of the API call.

(61) In some embodiments, the API authentication agent **228** uses modules (e.g., third-party credential authenticating programs) executing outside of the application **210** to authenticate the API call. In these embodiments, the authentication agent **228** passes a set of source-identifying parameters that identifies the source of the API call to the external module, so that this module can authenticate the source of the API call. One or more parameters in this set can be parameters that were received with the API call.

(62) After directing (at **315**) the authentication agent to authenticate the source of the API call, the API handler (at **320**) determines whether the authentication agent **228** reply indicates that the API call has been authenticated. If not, the API handler discards the API call and the process **300** ends.

In some embodiments, the API handler **226** not only discards the API call but also sends a reply to the source of the API call to reject the call.

(63) On the other hand, when the authentication agent **228** authenticates the source of the API call, the API handler **226** directs (at **325**) the application API-authorizing agent **230** to authorize the API call. In response, the application agent **230** sends an API-authorization request to the local API-authorizing agent **220**. To do this, application agent **230** generates (at **325**) an RPC message to the local agent **220** to pass along this request with information about the API call (e.g., one or more parameters that describe the API call and one or more parameters associated with this call). The network stack **215** passes this RPC message to the local agent **220**.

(64) After receiving the API-authorization request from application authorizing agent **230**, the local agent **220** uses (at **325**) one or more of the received parameters to retrieve a policy for authorizing the API call from the local storage **225**. In the embodiments in which the local storage **225** stores the policy and associated parameters in a namespace, the local agent uses one or more of the parameters in the API-authorization request to retrieve policy opcode related to the API call from the namespace. To evaluate this policy opcode, the agent might also retrieve parameters (e.g., policy operands) from the local storage (e.g., from the namespace) that are associated with the retrieved policy.

(65) The local agent then uses the received parameter set and/or retrieved parameter set to evaluate the retrieved API-authorization policy (e.g., the retrieved opcode) in order to determine whether the API call should be approved or rejected. In the embodiments that express each policy in terms of policy opcode that has one or more rules that specify one or more conditions for approving or rejecting the API calls, the agent resolves these conditions based on the received and/or retrieved parameters. For instance, in some embodiments, the agent determines that one or more parameters in the received parameter set and/or retrieved parameter set match a condition specified by the policy opcode for rejecting the API call. In some embodiments, the conditions can specify reasons for allowing an API call.

(66) Some API-authorization policies are resolved only by reference to the received parameter sets (i.e., only specify conditions that just need the received parameter sets to be resolved). Other API-authorization policies are resolved only by reference to the retrieved parameter sets (i.e., only specify conditions that just need the retrieved parameter sets to be resolved). Still other API-authorization policies are resolved by reference to both the received and retrieved parameter sets (i.e., specify conditions that need both the received and retrieved parameter sets to be resolved).

(67) After evaluating the retrieved policy, the local agent then sends (at **325**) a reply message to the application API-authorizing agent **230** to return the results of its evaluation (e.g., to specify whether the API call should be allowed or rejected). In some embodiments, this reply message is an RPC reply message that is directed back to the application agent **230** through the network stack **215**. Once the application agent **230** receives the reply from the local agent **220**, it passes the local agent's approval or rejection of the API call to the API handler.

(68) After directing (at **325**) the authorization agent to authorize the API call, the API handler (at **330**) determines whether the application API-authorizing agent's reply indicates that the API call has been authorized. If not, the API handler discards the API call and the process **300** ends. In some embodiments, the API handler **226** not only discards the API call but also sends a reply to the source of the API call to reject the call.

(69) On the other hand, when the authorization agent **230** returns a reply that indicates that the API call has been approved, the API handler **226** directs (at **340**) the API-processing module **224** that is associated with the API call to perform the operation associated with the API call and to provide a reply message to the source of the API call to indicate the completion of this operation. In some cases, this operation requires an output to be generated and returned to the source of the API call or to some other destination. In these cases, the API processing module generates this output and provides this output to the source of the API call or to the other destination. After **340**, the process

300 ends.

(70) One of ordinary skill will realize that the process **300** is implemented differently in other embodiments. For instance, the API-processing application does not perform the authorization/authentication check operations of the process **300**. Instead, a proxy authentication/authorization application performs these operations. In these embodiments, the proxy application first intercepts an API call intended for a second application, and sends a request (e.g., through RPC, IPC, etc.) to authorize this API call to the local agent. When the local agent authorizes this request, this proxy application then passes the API call to the second application, while discarding this call when the local agent rejects this call. In some embodiments, the proxy application not only intercepts API calls to obtain authorization for these calls, but also intercepts these calls to first have them authenticated by a separate authentication service.

(71) Also, the process **300** discards an API call when this call is not authenticated. In other embodiments, the process **300** does not discard such an API call. For instance, in some embodiments, the process **300** passes the authentication failure of an API call to the local agent. This agent then uses this failure like a received parameter to resolve one or more conditions associated with the policy (retrieved from its namespace portion) that it evaluates to determine whether to authorize the API call.

(72) FIG. 4 illustrates a more-detailed view of a local agent **420** of some embodiments. In this example, the local agent **420** uses a local namespace portion **425** as the local policy and data storage **225**. This namespace contains both (1) policy opcodes to evaluate API calls to determine whether they should be authorized, and (2) operands to use to evaluate and resolve the opcodes. In some embodiments, the namespace portion **425** contains the policy opcodes and operands needed for the API calls that are received by one or more applications on the agent's host that use the local agent. In other embodiments, the local agent has a different namespace portion for different applications.

(73) The local agent **420** has a namespace interface **422** that receives and stores the namespace portion and updates to all or part of this portion from the remote policy server set **105**. This interface also receives policy operands from local parameter integrator **255**, which, as mentioned above, can be used in some deployments to store policy operands in the namespace that are not provided to the server set **105** (e.g., to maintain the confidentiality of such policy operands). Upon the initialization of the local agent **420**, the namespace interface **422** in some embodiments initially queries the server set to obtain the namespace(s) for the applications that will use the local agent. Also, in some embodiments, the namespace interface **422** periodically polls the server set to determine whether any policy or operand in its namespace has been added, deleted or modified, and if so, the namespace interface updates its local namespace portion **425** to reflect this change (e.g., to download any new and/or modified policy and/or operand and store the downloaded policy and/or operand in the namespace portion).

(74) In addition to the namespace interface **422**, the local agent **420** includes a request handler **405**, evaluation engine **410**, policy and data fetcher **415**, and rule compiler **430**. The operations of these modules will be described by reference to FIG. 5, which illustrates a process **500** that the local agent **420** performs when it receives a request to authorize an API call from an API-authorizing agent **230** of an application **210** that executes on the same host as the local agent **420**.

(75) As shown, the process **500** starts (at **505**) when the agent's API-authorization request handler **405** receives an API-authorization request from the application's agent **230**. This request is received in some embodiments as an RPC message that includes a set of parameters associated with the API call that the application has received. In some embodiments, this parameter set includes a parameter that identifies the API call and one or more metadata parameters associated with the API call. Also, in some embodiments, this parameter set includes one or more parameters that the local agent needs to resolve policy opcode that is associated with the API call. Accordingly, after receiving the RPC message, the handler extracts (at **510**) this parameter set from the RPC

message's payload and provides this parameter set to the evaluation engine **410**. This parameter set identifies the API call as mentioned above.

(76) The evaluation engine is responsible for evaluating policy opcode related to the API call (i.e., the API call specified in the API-authorization RPC message) based on one or more parameters received with the RPC message and/or retrieved from the local namespace. Before evaluating this opcode based on one or more associated parameters, the evaluation engine needs the policy and data fetcher **415** to retrieve from the local namespace portion **425** the policy opcode and operands applicable to the identified API call, and the rule compiler **430** to create a more optimal runtime rule and parameter structure **435** for the evaluation engine to process. The evaluation engine also stores compiled rule structures for prior API-authorization requests in a cache storage **440**, so that it can forego retrieving and compiling the same policy opcode/operands when it subsequently receives the same API call.

(77) In view of this functionality, after identifying (at **510**) the API call associated with the received request, the evaluation engine determines (at **515**) whether the cache storage **440** stores a reference to a rule structure for this API call that was previously compiled when the local agent received the same API-authorization request previously. In some embodiments, the cache **440** identifies each stored rule structure in terms of one or more parameters that are passed along with the API-authorization request to the local agent **420** by the API-authorization agent **230**. Hence, in these embodiments, the evaluation engine determines whether the parameters passed along with the current request match the parameters stored for the rule structures identified in the cache storage **440**.

(78) If so, the evaluation engine processes (at **520**) this previously specified rule structure to formulate its decision, and then provides (at **520**) its decision (i.e., its “authorization” or “rejection”) to the request handler **405**. At **520**, the API handler then formulates an RPC reply message with this decision in its payload and sends this reply message to the authorization agent **230** that sent the request. After **520**, the process **500** ends.

(79) On the other hand, when the evaluation engine determines (at **515**) that the cache storage does not store a reference to a previously defined rule structure for the API call identified in the received request, the evaluation engine directs (at **525**) the policy and data fetcher **415** to retrieve from the local namespace portion **425** the policy opcode and operands applicable to the identified API call. As described above and further described below, the namespace is a hierarchical structure that stores policy opcode associated with an API call in a location (e.g., a path) in the structure that can be specifically identified based on the API call's identifier (e.g., the name associated with the API call). In some embodiments, this location also stores the operands needed for processing this policy opcode and/or stores a location identifier (e.g., the path) that specifies other location(s) in the namespace that store the operands needed for processing the policy opcode. Hence, in some embodiments, the fetcher **415** can easily identify the policy opcode and operands to retrieve by just using parameters associated with the API call.

(80) After retrieving the policy opcode and one or more associated operands, the fetcher **415** directs (at **530**) the rule compiler **430** to create a more optimal runtime rule and parameter structure **435** for the evaluation engine to process. As described above, the policy opcode for an API call includes one or more rules, with each rule expressing one or more conditions for rejecting (or allowing) the API call. In some embodiments, the optimal runtime rule and parameter **435** structure includes “prepared rules” that are generated by the compiler. Prepared rules are rules that are parsed from the retrieved, compiled and optimized policy opcode that is retrieved by the fetcher for the API call.

(81) Prepared rules are ready for execution and thus require less effort to execute. In some embodiments, the prepared rules are expressed using an abstract syntax tree (AST). Some embodiments translate the AST into a sequence of executable instructions. These instructions are virtual instructions (not CPU instructions) that are easy to run fast, and even further optimize for speed. Some embodiments use known processes to compile and process prepared rules from

declarative language code. Some of the known processes are described in Handbook of Automated Reasoning by Alan Robinson and Andrei Voronkov.

(82) Once the rule compiler creates (at **530**) the more optimal rule structure, it notifies (at **535**) the evaluation engine **410** directly or through the fetcher **415**. The evaluation engine then processes (at **535**) the rules in this rule structure. In some embodiments, the evaluation engine uses the parameters retrieved from the namespace and/or received with the API request to resolve these rules (e.g., conditions specified by the rules).

(83) After processing (at **535**) the rule structure, the evaluation engine **410** provides its decision with respect to the API call (i.e., the “authorization” or “rejection” of the API call) to the handler **405**. At **540**, the handler then formulates an RPC reply message with this decision in its payload and sends this reply message to the authorization agent **230** that sent the request. Next, at **545**, the evaluation engine stores in the cache **440** a reference to the optimized rule structure that it created at **530**. This cached result can be kept as long as the retrieved policy opcode does not change: that way the evaluation engine can skip the parsing/compiling/optimization part (which can be more time consuming) and quickly start executing with minimal effort. As mentioned above, the evaluation engine **410** stores the cached rule structure by using a reference that is associated with one or more parameters that were received with the API call or the API-authorization request. After **545**, the process ends.

(84) In some embodiments, the process **500** not only caches the optimized rule structure for an API call's retrieved policy, but also caches the decisions that it reaches in processing the rule structure for a particular set of parameters. However, as each decision is dependent on the set of parameters that can dynamically change or can resolve the rules differently at different times or under different conditions, the process has to ensure that the previously cached decisions are still valid at a later instance in time when the API call is received again.

(85) As mentioned above, the local agents in some embodiments are used to authorize API calls for both application domain logic and application infrastructure logic, while in other embodiments, the local agents are used to authorize API calls for application infrastructure logic but not application domain logic. Also, as mentioned above, decoupling API authorization for infrastructure logic simplifies the application code and allows the application's code to focus on the matters relevant to the application's core functions. This is especially the case because while infrastructure functions are critical, they are not typically a concern of the application domain logic.

(86) For instance, in banking, it is the infrastructure part of the application that is responsible for maintaining micro-services' connections to an account database holding the actual account data (there could many of those), providing the necessary HTTP facilities to allow the customers' browsers to retrieve their account contents (customers would expect SSL, for instance), as well as providing connectivity between micro-services (modern applications include several micro-services which then together form the application).

(87) Both core logic and infrastructure parts require authorization in their operations. However, the type of authorization implemented is different: core logic implements authorization decisions over the core data and operations, while the infrastructure parts implement authorization decisions for the functionality they provide.

(88) This implies that core logic is likely to implement authorization decisions that are fundamentally tied to the logic being implemented. For instance, in banking, there are restrictions about the mortgage offer approvals, between which account transfers can be moved without further security inspection (bank internal transfers are less risky compared to transfers to other banks), or between which kind of accounts fund transfers are allowed to begin with (a savings account could have restrictions for the number of transfers done per month). Even the size of the transfer itself is likely to affect the authorization decision.

(89) Similarly in healthcare, the core logic is the one that is responsible for health care data and operations over it: creation and filling prescription is likely to involve several authorization steps

that are tied to the processing of the prescription data. For instance, heavy painkillers are likely to have more strict authorization logic in place to allow the creation of the prescription to begin with. (90) On the contrary, authorization within the infrastructure parts of the application and micro-service revolve around the infrastructure level entities. The authorization decisions are about whether a micro-service can talk to another, which micro-service API calls can invoke, or whether a micro-service can open a connection to a particular database instance and read a particular table. For business logic, these are low-level details that are typically invisible to the end user.

(91) Application infrastructure level authorization is also relevant and related to the managing of the application and micro-service as whole. For example, a micro-service instance might run inside a VM and the type of a micro-service is likely to affect who (and when) can access the VM. For instance, it is common to block all access by default and allow access only as-needed basis. In that case only the people actively resolving an incident related to application are allowed to access the VM and hence the application deployment. While this does not concern the application instance itself (a process running a micro-service instance is not responsible for managing the VM the process runs within), these authorization decisions can be considered to be part of application infrastructure authorization.

(92) As mentioned above, the server set **105** in some embodiments collects and updates parameters for resolving policies dynamically and in a deployment specific manner (i.e., in a manner that can be context specific for each set of associated machines for which the server set is deployed). This allows policies to be resolved based on conditions and dynamically changing parameters that are specific to each deployment. As further described below, the server set allows the values for policy operands to dynamically change because it allows local agents to receive (through push operations or pull operations) such changes for the policy operands in the respective namespaces of the local agents.

(93) FIGS. **6-8** present three examples that illustrate the evaluation engine **410** of a local agent resolving the same API authorization request based on deployment specific policies and operands, with the values of the operands changing dynamically based on network conditions. In these examples, the same application (not shown) sends three requests to the evaluation engine to authorize an API call associated with changing a setting of a VM. Also, in each of these examples, the evaluation engine processes the same API-authorization policy, which includes one rule that specifies that modifications to the VM settings should be rejected during time periods in which an incident is being reported for the VM's host computer unless the user requesting the setting change has been issued a ticket, such as a user-authentication ticket or a project tracking ticket. In some embodiments, the ticket would be issued by a third-party ticketing application, such the Jira Issue and Project Tracking applications of Atlassian.

(94) In these examples of FIGS. **6-8**, the policy being enforced is custom-defined for a deployment of the policy enforcement system **100** of some embodiments. While this deployment restricts VM modifications during incidents, another deployment of the policy enforcement system might not. Other deployments might not even have any VMs; they might only have standalone computers or they might only use containers. Moreover, in these examples, the operands for resolving the specified VM-modification policy dynamically change. Specifically, these operands include (1) the Boolean Incident flag and (2) array that specifies identity of the users with tickets (if any). Both the Boolean parameter and the ticket array can be changed at any time. The Incident flag is false when there is no issue with the VM's host computer, while it is true that when there is an issue with the VM's host computer. Also, each time a user obtains a ticket for performing a task, the user's identity is added to the ticket array. The user's identity is removed from the ticket array when the ticket expires. All of these parameters are deployment specific and can change dynamically in real time.

(95) As mentioned above, the agent's local namespace portion **425** in some embodiments is updated (through push or pull operations) whenever policy opcodes and/or policy operands in the copy of the namespace that is maintained by the server set **105**. Similarly, the server set's copy of the

namespace is updated whenever a new policy is added, modified or deleted, and/or a operand's value changes through the various policy input interfaces and operand import interfaces of the server set.

(96) Accordingly, whenever the Incident flag changes its value or whenever a user is added or removed from the ticket array, the server set's operand import interfaces modify the values of the Incident flag and the ticket array. From this namespace, these changes are propagated to the agent's local namespace portion **425** through push operations or pull operations. The examples illustrated in FIGS. **6-8** provide an abridged view of this propagation by showing data being sent from one or more data sources **605** in the deployment environment to the local namespace portion **425**. These figures also do not show the policy and data fetcher **415**, the rule compiler **430** and the created runtime rule and parameter structure **435** in order to keep the illustrations in these figures simple. Instead, these figures just show the evaluation engine **410** receiving the VM-modification policy's rule that specifies that VM changes should be rejected during an incident when the requesting user does not have a ticket. These abridged representations are illustrated with dashed arrows in these figures.

(97) FIG. **6** illustrates the case when the VM modification API call is received when there is no incident. Thus, in this example, the evaluation engine sends a reply **610** that indicates that the API call is authorized. FIG. **7** illustrates the case when the VM modification API call is received when there is an incident and the user associated with this API call is Alice. As shown, the ticket array **615** identifies Alice as someone who has been issued a ticket. Accordingly, in this example, the evaluation engine sends a reply **620** that indicates that the API call is authorized. FIG. **8** illustrates the case when the VM modification API call is received when there is an incident but the user associated with this API call is Bob. As shown, the ticket array **625** does not identify Bob as someone who has been issued a ticket; it only identifies Alice and Jack as people who have been issued tickets. Accordingly, in this example, the evaluation engine sends a reply **630** that indicates that the API call is rejected.

(98) FIG. **9** illustrates an example of the policy opcode and operand that can be used to implement API-authorization controls similar to those of the examples illustrated in FIG. **6-8**. Unlike the examples of FIGS. **6-8** that resolve one rule, the policy opcode in the example of FIG. **9** has three rules. To resolve this policy opcode, three sets of parameters **902** are identified in-line within the policy opcode for authorizing API calls regarding VM modifications. These three sets of parameters specify an LDAP group directory, a data store specifying occurrences of incidents, and a JIRA ticketing data source.

(99) A first rule **905** specifies that a VM-modification API call should be allowed if the call comes from a user in the engineering staff during a period in which VM-modifications are not blocked. The second rule **910** defines whether the current time period (i.e., the period during which the VM-modification API call is received) is a time period that the VM-modification call should be blocked. This rule specifies that the current time period is a blocked time period when the Boolean Incident flag is currently true and there is no temporary exemption. The third rule **915** defines whether an exemption should be granted during an incident for the user associated with the received API call. This rule specifies that a temporary exemption should be granted when the ticket array includes the user identity and the issues label for this user in the labels array identifies the ticket to be associated with the existence of an incident.

(100) Based on these rules, the evaluation engine will authorize VM-modification API calls from anyone within the engineering staff (as identified by the LDAP directory) when there is no incident. It will also authorize such calls during an incident from anyone within the engineering staff when the user has a ticket associated with incidents. However, it will reject such calls during an incident from anyone that does not have a ticket associated with the incident.

(101) The above-described local API-authorization request processing of FIGS. **2-9** has several advantages. First, by packing the relevant portion of the namespace with both policies and data

parameters for resolving these policies, and then providing this portion to a local agent executing on the same host computer as the applications that need their API calls authorized, this API-authorization request processing is fast and highly available. No remote device has to be accessed to process these requests.

(102) This approach can also cache the optimized rule data structure for the API-authorization policy opcode associated with an API call after this structure is created for assessing whether the API call should be allowed. This cached rule data structure can be subsequently use for processing requests associated with the same API call. In this manner, the cached structure eliminates the need for the rule compiler to recreate the optimized rule structure while the local agent processes the authorization request with the subsequent iterations of the same API call.

(103) As mentioned above, the local API-authorization request processing of some embodiments also caches decision. Whenever decisions are cached, the system has to address cache revocation. Revoking cache decisions when policies are enforced by a central set of servers is complicated. A centralized approach might use rather long caching timeouts to maximize the efficiency of its decision caching. But this would then require some mechanism to explicitly detect which cached decisions have to be revoked when a policy or data changes. Otherwise the local cache could store stale decisions for long periods of time, before they timeout and replaced with new ones matching the updated policy and data.

(104) To figure out the decisions to revoke from cache, the application in the centralized approach would have to keep track of policies and data used in the decisions it receives from the centralized location—and then the centralized decision maker would have to inform the application if any of these changes. This is convoluted at best: to support revocation, the interface between application and centralized decision maker has to deal with policy and data information that they would not otherwise even need. After all, the authorization decisions are executed centrally, outside of the application, yet, somehow applications would have to understand the data and policies used in those decisions so that they could cache those decisions. As a result, the interface between the application and central location would now include information about data, policies and decisions.

(105) On the other hand, in the above-described local API-authorization request processing of FIGS. 2-9, the interface between the application and decision is just about namespace (i.e., about the policies and data). The entire namespace bundle comes with a single version that the application can use to check if the central location has any updates for this portion of the namespace. This results in much simpler and reliable interface between the application and central location.

(106) One of ordinary skill will realize that FIGS. 6-9 only provide some examples that describe the operations of some embodiments of the invention. The API-authorization system of some embodiments can be used to resolve many other API-authorization policies dynamically. Moreover, even the operations described in the examples of FIGS. 6-9 are implemented differently in different embodiments. For instance, instead of relying on an incident flag to determine whether there is an incident associated with a host computer, VM or container on which an API-processing application executes, the evaluation engine 410 dynamically computes this parameter (i.e., dynamically determines whether there is an incident). To dynamically determine this parameter, the evaluation engine 410 in some embodiments could direct the fetcher 415 to iteratively retrieve different policies and/or operands from the namespace as the evaluation engine processes the optimized rule structure 435. The incident parameter could then be the result of one or more of these retrieval and processing operations.

(107) FIG. 10 illustrates the software architecture of the server set 105 of some embodiments of the invention. Two or more servers in some embodiments can implement the server-side modules 1010 in this architecture. Also, in some embodiments, multiple instances of the same modules execute on different servers in the server set 105 as this server set has two or more servers performing the same operations in these embodiments. As shown, the server set 105 includes local-agent API

interface **1025**, browser-API interface **1030**, policy and content document database **1035**, indexed decision logs and agent status data storage **1042**, parameter ingesting interface **1040** and policy authoring interface **1045**.

(108) The interfaces **1025** and **1030** communicatively connect the server-side modules **1010** to client-side modules **1005**. Specifically, the browser interface **1030** processes communications between the server set's modules and the browser clients, while the agent interface **1025** processes communications between the server set modules and the remote local agents **220** on the host computers.

(109) The local agents **220** communicate with server modules **1010** by using APIs of the server's agent interface **1025**. This interface includes (1) a download API that provides the functionality to download a bundle of policies and related data, and then request its next version, when a new version is available, and (2) an upload API that provides the functionality for the agents to report back any policy decisions they have made. The download API accesses the document database underneath to compile the policies and any of their dependencies (whether they are data or other policies) into a bundle. In some embodiments, the agent interface **1025** provides local namespace portions and updates to these portions in response to queries for the namespaces and updates to the namespaces. The upload API uses a separate decision database to log any decisions clients have reported.

(110) As mentioned above, the policy and content document database **1035** in some embodiments is a custom namespace database. This database can persist both policies and any data the policies need in their evaluation. In some embodiments, the database saves policies as plain source code, while storing the parameter data as JSON documents. The policy and content database **1035** in some embodiments includes a hierarchical namespace for each entity in one or more datacenters for which the server set defines and distributes API-authorization policies. As this database **1035** merges the parameter data and policies into one global hierarchical namespace layout, database clients can access a nested portion of the namespaces simply by requesting policy operands by their fully qualified names.

(111) The server set copy of the policy and content database **1035** in some embodiments can save all versions of the parameter data and policies over time for an entity. When either a data or policy is updated, its previous version is not removed but archived in some embodiments. This allows time traversal: the server set can inspect the state of policies and their related parameter data at any point in time by simply defining the time of interest when accessing the namespace. Time traversal is useful for troubleshooting (to understand exactly what was the state of the policies at the time of an issue) and for auditing (to convince the administrator that the system state was correct throughout the audited period of time). Also, in some embodiments, the policy and content database **1035** implements delta-encoding to avoid saving entire copies of all data and policies after each, potentially small, change. This is very effective as most of the data is rather stable with only a steady stream of fairly local changes over time.

(112) As mentioned above, policies refer and use deployment-specific parameters describing the current state of the managed resources and datacenter. To retrieve and properly format the parameters for consumption, the server set uses the parameter ingesting interface **1040** in some embodiments. The parameter ingesting interface **1040** includes one or more data source adapters, each capable of using a remote management API of a parameter data source to collect deployment specific parameters (e.g., through data crawl operations) and transforming the collected parameter data from its native format into a structured document (e.g., a JSON document) for storing in the database **1035**. For instance, when the API authorization system validates API requests to micro service applications deployed in Amazon Web Services (AWS) cloud infrastructure, the parameter ingestion interface **1035** has one or more data source adapters to collect data needed for processing different AWS APIs.

(113) In case a relevant data is easily accessible from a database, the parameter-ingesting interface

1040 includes a data source adapter that can implement the relevant database access protocol and directly replicate the relevant portion of the remote database. For instance, in some embodiments, a data source adapter is used to retrieve user group information from a remote user directory by using LDAP protocols.

(114) The policy-authoring interface **1045** in some embodiments provides resources that allow the server set to generate policies for an entity. In some embodiments, this interface is accessible through browser clients **1020** and the browser interface **1030**. Policies in some embodiments can also be uploaded to the server set through an API interface.

(115) In some embodiments, the decision and status storage **1042** is a database that stores the received decision log entries for archiving. An indexing engine (not shown) builds a text search index over the stored decision log entries. The index is useful for the UI component as it enables running queries about the decisions, e.g., during troubleshooting when a particular application is not operating as expected and the related policy decisions need to be further inspected. The browser clients **1020** in some embodiments are JavaScript applications. Through the browser API interface **1030**, the browser clients **1020** can access the underlying databases/indices. This allows the browser clients **1020** to access any current or past policy/parameter state, as well as allow inspection of the policy decisions done.

(116) The policy defining and distribution system **1000** of FIG. **10** has several advantages. It imports remote management API state (e.g., as JSON objects) in the document database, which allows the system to effectively construct a logical, centralized view of the datacenter state. This combined with versioning allows the system to provide versioned snapshots of the entire datacenter state, if needed. For example, in some embodiments, the local agents **220** identify the namespace versions that they use to process the API-authorization requests and provide this version information to the server set to store in the decision logs **1042**. Subsequently, through the browser interface **1030** and the decision logs **1042**, a user can perform browser-based searches to extract the exact state of the policies and parameters at the time a decision was made. This state information includes the version data. The version data can also be used to search the decision logs. These capabilities are highly useful for debugging the API-authorization system. These search and debug operations are also supported in some embodiments through API interfaces that support API calls from automated processes that collect and process log data.

(117) Also, the system's overall process of extracting remote management APIs, disseminates relevant pieces of this data together with the relevant policy code for the agents. This is near real time dissemination of management state is highly beneficial that current techniques cannot duplicate. Current systems integrate with the necessary remote management APIs on their own, which creates N.sup.2 integrations, whereas the system **100** can do the integrations once, which results N integration.

(118) The system **100** also uses a declarative policy language for document-oriented, structured data. Also, as mentioned above, the system **100** in some embodiments decouples policy (authorization) logic from the applications to an agent nearby. In a traditional authorization solution (e.g., using LDAP), the policy logic is part of the application source code and application directly retrieves the relevant data from the LDAP, applying its logic on the data retrieved.

(119) Many of the above-described features and applications are implemented as software processes that are specified as a set of instructions recorded on a computer readable storage medium (also referred to as computer readable medium). When these instructions are executed by one or more processing unit(s) (e.g., one or more processors, cores of processors, or other processing units), they cause the processing unit(s) to perform the actions indicated in the instructions. Examples of computer readable media include, but are not limited to, CD-ROMs, flash drives, RAM chips, hard drives, EPROMs, etc. The computer readable media does not include carrier waves and electronic signals passing wirelessly or over wired connections.

(120) In this specification, the term “software” is meant to include firmware residing in read-only

memory or applications stored in magnetic storage, which can be read into memory for processing by a processor. Also, in some embodiments, multiple software inventions can be implemented as sub-parts of a larger program while remaining distinct software inventions. In some embodiments, multiple software inventions can also be implemented as separate programs. Finally, any combination of separate programs that together implement a software invention described here is within the scope of the invention. In some embodiments, the software programs, when installed to operate on one or more electronic systems, define one or more specific machine implementations that execute and perform the operations of the software programs.

(121) FIG. 11 conceptually illustrates an electronic system **1100** with which some embodiments of the invention are implemented. The electronic system **1100** may be a computer (e.g., a desktop computer, personal computer, tablet computer, server computer, mainframe, a blade computer etc.), phone, PDA, or any other sort of electronic device. The electronic system **1100** is also the control plane modules of the load balancer of some embodiments. As shown, the electronic system includes various types of computer readable media and interfaces for various other types of computer readable media. Specifically, the electronic system **1100** includes a bus **1105**, processing unit(s) **1110**, a system memory **1125**, a read-only memory **1130**, a permanent storage device **1135**, input devices **1140**, and output devices **1145**.

(122) The bus **1105** collectively represents all system, peripheral, and chipset buses that communicatively connect the numerous internal devices of the electronic system **1100**. For instance, the bus **1105** communicatively connects the processing unit(s) **1110** with the read-only memory (ROM) **1130**, the system memory **1125**, and the permanent storage device **1135**. From these various memory units, the processing unit(s) **1110** retrieve instructions to execute and data to process in order to execute the processes of the invention. The processing unit(s) may be a single processor or a multi-core processor in different embodiments.

(123) The ROM **1130** stores static data and instructions that are needed by the processing unit(s) **1110** and other modules of the electronic system. The permanent storage device **1135**, on the other hand, is a read-and-write memory device. This device is a non-volatile memory unit that stores instructions and data even when the electronic system **1100** is off. Some embodiments of the invention use a mass-storage device (such as a magnetic or optical disk and its corresponding disk drive) as the permanent storage device **1135**.

(124) Other embodiments use a removable storage device (such as a floppy disk, flash drive, etc.) as the permanent storage device. Like the permanent storage device **1135**, the system memory **1125** is a read-and-write memory device. However, unlike storage device **1135**, the system memory is a volatile read-and-write memory, such as a random access memory. The system memory stores some of the instructions and data that the processor needs at runtime. In some embodiments, the invention's processes are stored in the system memory **1125**, the permanent storage device **1135**, and/or the read-only memory **1130**. From these various memory units, the processing unit(s) **1110** retrieve instructions to execute and data to process in order to execute the processes of some embodiments.

(125) The bus **1105** also connects to the input and output devices **1140** and **1145**. The input devices enable the user to communicate information and select commands to the electronic system. The input devices **1140** include alphanumeric keyboards and pointing devices (also called “cursor control devices”). The output devices **1145** display images generated by the electronic system. The output devices include printers and display devices, such as cathode ray tubes (CRT) or liquid crystal displays (LCD). Some embodiments include devices such as a touchscreen that function as both input and output devices.

(126) Finally, as shown in FIG. 11, bus **1105** also couples electronic system **1100** to a network **1165** through a network adapter (not shown). In this manner, the computer can be a part of a network of computers (such as a local area network (“LAN”), a wide area network (“WAN”), or an Intranet, or a network of networks, such as the Internet. Any or all components of electronic system **1100** may

be used in conjunction with the invention.

(127) Some embodiments include electronic components, such as microprocessors, storage and memory that store computer program instructions in a machine-readable or computer-readable medium (alternatively referred to as computer-readable storage media, machine-readable media, or machine-readable storage media). Some examples of such computer-readable media include RAM, ROM, read-only compact discs (CD-ROM), recordable compact discs (CD-R), rewritable compact discs (CD-RW), read-only digital versatile discs (e.g., DVD-ROM, dual-layer DVD-ROM), a variety of recordable/rewritable DVDs (e.g., DVD-RAM, DVD-RW, DVD+RW, etc.), flash memory (e.g., SD cards, mini-SD cards, micro-SD cards, etc.), magnetic and/or solid state hard drives, read-only and recordable Blu-Ray® discs, ultra density optical discs, any other optical or magnetic media, and floppy disks. The computer-readable media may store a computer program that is executable by at least one processing unit and includes sets of instructions for performing various operations. Examples of computer programs or computer code include machine code, such as is produced by a compiler, and files including higher-level code that are executed by a computer, an electronic component, or a microprocessor using an interpreter.

(128) While the above discussion primarily refers to microprocessor or multi-core processors that execute software, some embodiments are performed by one or more integrated circuits, such as application specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs). In some embodiments, such integrated circuits execute instructions that are stored on the circuit itself.

(129) As used in this specification, the terms “computer”, “server”, “processor”, and “memory” all refer to electronic or other technological devices. These terms exclude people or groups of people. For the purposes of the specification, the terms display or displaying means displaying on an electronic device. As used in this specification, the terms “computer readable medium,” “computer readable media,” and “machine readable medium” are entirely restricted to tangible, physical objects that store information in a form that is readable by a computer. These terms exclude any wireless signals, wired download signals, and any other ephemeral or transitory signals.

(130) While the invention has been described with reference to numerous specific details, one of ordinary skill in the art will recognize that the invention can be embodied in other specific forms without departing from the spirit of the invention. For instance, a number of the figures conceptually illustrate processes. The specific operations of these processes may not be performed in the exact order shown and described. The specific operations may not be performed in one continuous series of operations, and different specific operations may be performed in different embodiments. Furthermore, the process could be implemented using several sub-processes, or as part of a larger macro process.

(131) While micro-services were mentioned as one type of applications that could get the benefit of the API-authorization processes and systems of some embodiments, one of ordinary skill will realize that these processes and systems of some embodiments are also used to authorize API calls for other types of applications, such as SSH (secure shell) applications or other types of applications. Also, in several embodiments described above, the API-authorization enforcement system enforces API-authorization policies for different sets of related machines in one or more datacenters. In other embodiments, the API-authorization enforcement system enforces API-authorization policies just based on the applications or types of applications executing on the computers. For instance, a datacenter can employ the API-authorization enforcement system of some embodiments for one or more application or application types of all of its tenants or users. In these embodiments, the policies are policies that are deployed based on application identity or type, and not based on association between related sets of machines executing in the datacenter.

(132) Also, several of the above-described embodiments have the local agents and/or the server set evaluate the policies based on the operands to determine whether to allow or reject API calls. In other embodiments, the local agents or the server set configure the applications **115** that receive the API calls, or other agents/processes in the datacenter authorize or reject API calls based on

configurations that the local agents and/or server set push to these applications and/or other agents/processes. For instance, each time a local namespace portion is updated, the local agent or the server set in an AWS datacenter identifies the API calls affected by the update, determines whether these API calls should now be allowed or rejected, and pushes to the AWS IAM (Identity and Access Management) system new configuration data to ensure that these API calls are appropriately allowed or rejected.

(133) In some embodiments, some or all of the local agent functionality is embedded in a plugin that is installed in the API-processing application. In this manner, some or all of the above-described operations of the local agents are performed by plugins installed in the API-processing applications in some embodiments. In other embodiments, instead of implementing these operations with plugins, some embodiments have the local agent and/or server set update API rule configuration file of an API-processing application whenever the local namespace associated with the application is modified and this modification affects the application's processing of one or more API calls. Therefore, one of ordinary skill in the art would understand that the invention is not to be limited by the foregoing illustrative details, but rather is to be defined by the appended claims.

Claims

1. A system for authorizing API (Application Programming Interface) calls, the system comprising: a plurality of computers, each computer executing (i) at least one application that receives API calls and (ii) at least one local API-authorizing agent to authorize API calls received by the at least one application executing on the same computer; and a set of servers to receive API-authorization policies and parameters for evaluating the policies and to distribute the policies and parameters to the local agents to use to authorize the API calls received by applications executing on the same computers as the local agents, wherein each respective application of a plurality of applications executing on a particular computer processes API calls received by the respective application by (i) directing the received API calls to a respective local API-authorizing agent executing on the same particular computer as the respective application for the respective local API-authorizing agent to authorize the API calls for processing by the respective application, (ii) receiving authorization decisions for the API calls from the respective local API-authorizing agent, and (iii) for each API call authorized by the respective local API-authorizing agent, performing an operation specified by the API call at the application and sending a message to a source of the API call to indicate that the operation was performed, wherein the local API-authorizing agents executing on the particular computer authorize the API calls by (i) identifying, for each particular API call of a plurality of the API calls, an API-authorization policy applicable to the particular API call and (ii) evaluating the identified API-authorization policy for the particular API call, wherein a particular local API-authorizing agent on the particular computer (i) uses a first set of parameters associated with a particular API call to retrieve, from a single storage structure that stores both API-authorization policies and parameters for evaluating policies, at least one API-authorization policy and a second set of parameters for evaluating the retrieved API-authorization policy, and (ii) uses the retrieved second set of parameters to evaluate the retrieved API-authorization policy in order to determine whether the particular API call should be approved, wherein the particular local API-authorizing agent receives the single storage structure from the set of servers.

2. The API-authorizing system of claim 1, wherein the parameters in the storage structure include at least one parameter received from a local data collector that executes on the same particular computer as the particular local API-authorizing agent, said parameter received from the local data collector comprising confidential data not provided to the server set.

3. The API-authorizing system of claim 1, wherein the single storage structure is a hierarchical storage structure that allows each respective API-authorization policy and a respective set of parameters associated with the respective policy to be retrieved by providing a respective location

in the hierarchical storage structure that stores the respective policy.

4. The API-authorizing system of claim 3, wherein the storage structure is a namespace.

5. The API-authorizing system of claim 3, wherein: the hierarchical storage structure is a namespace; and the location that stores a particular policy also stores one of (i) the set of parameters for evaluating the particular policy and (ii) a name of a location in the namespace at which the set of parameters for evaluating the particular policy has been previously stored in the namespace.

6. The API-authorizing system of claim 1, wherein: each computer further executes at least one authenticating agent to authenticate the API calls received by the at least one application executing on the same computer; and each respective application executing on the particular computer further processes received API calls by directing the API calls to a respective authenticating agent for authentication prior to directing the API calls to the respective API-authorizing agent for authorization.

7. The API-authorizing system of claim 1, wherein the set of servers further authorizes API calls for a subset of the applications executing on the plurality of computers.

8. The API-authorizing system of claim 7, wherein a subset of the local API-authorizing agents relay API authorization requests from the subset of the applications to the set of servers.

9. The API-authorizing system of claim 1, wherein the received API calls are for (i) a first set of applications that provide infrastructure logic and (ii) a second set of applications that provide domain logic that specifies domain rules for creating, storing, and manipulating data relating to core functions of an enterprise.

10. The API-authorizing system of claim 1, wherein the set of servers executes a set of data-source adapters to collect a subset of the parameters.

11. A system for authorizing API (Application Programming Interface) calls, the system comprising: a plurality of computers, each computer executing (i) at least one application that receives API calls and (ii) at least one local API-authorizing agent to authorize API calls received by the at least one application executing on the same computer; and a set of servers to receive API-authorization policies and parameters for evaluating the policies and to distribute the policies and parameters to the local agents to use to authorize the API calls received by applications executing on the same computers as the local agents, wherein each respective application of a plurality of applications executing on a particular computer processes API calls received by the respective application by (i) directing the received API calls to a respective local API-authorizing agent executing on the same particular computer as the respective application for the respective local API-authorizing agent to authorize the API calls for processing by the respective application, (ii) receiving authorization decisions for the API calls from the respective local API-authorizing agent, and (iii) for each API call authorized by the respective local API-authorizing agent, performing an operation specified by the API call at the application and sending a message to a source of the API call to indicate that the operation was performed, wherein the local API-authorizing agents executing on the particular computer authorize the API calls by (i) identifying, for each particular API call of a plurality of the API calls, an API-authorization policy applicable to the particular API call and (ii) evaluating the identified API-authorization policy for the particular API call, wherein the set of servers receives records expressing the API authorization decisions made by the local agents and stores the received records in a set of data storages for querying to examine API authorizations decisions made by the local agents.

12. The API-authorizing system of claim 11, wherein a particular local API-authorizing agent on the particular computer (i) uses a first set of parameters associated with a particular API call to retrieve, from a single storage structure that stores both API-authorization policies and parameters for evaluating policies, at least one API-authorization policy and a second set of parameters for evaluating the retrieved API-authorization policy, and (ii) uses the retrieved second set of parameters to evaluate the retrieved API-authorization policy in order to determine whether the

particular API call should be approved.

13. The API-authorizing system of claim 12, wherein the particular local API-authorizing agent receives the single storage structure from the set of servers.

14. A system for authorizing API (Application Programming Interface) calls, the system comprising: a plurality of computers, each computer executing (i) at least one application that receives API calls and (ii) at least one local API-authorizing agent to authorize API calls received by the application executing on the same computer; and a set of servers to receive API-authorization policies and parameters for evaluating the policies and to distribute the policies and parameters to the local agents to use to authorize the API calls received by applications executing on the same computers as the local agents, wherein each respective application of a plurality of applications executing on a particular computer processes API calls received by the respective application by (i) directing the received API calls to a respective local API-authorizing agent executing on the same particular computer as the respective application for the respective local API-authorizing agent to authorize the API calls for processing by the respective application and (ii) receiving authorization decisions for the API calls from the respective local API-authorizing agent, wherein at least a first application executing on the particular computer directs received API calls to a first local API-authorizing agent executing on the particular computer and a second application executing on the particular computer directs received API calls to a second local API-authorizing agent executing on the particular computer, wherein the first local API-authorizing agent authorizes API calls made to the first application and a first plurality of additional applications executing on the particular computer while the second local API-authorizing agent authorizes API calls made to the second application and a second plurality of additional applications executing on the particular computer, wherein the local API-authorizing agents executing on the particular computer authorize the API calls by (i) identifying, for each particular API call of a plurality of the API calls, an API-authorization policy applicable to the particular API call and (ii) evaluating the identified API-authorization policy for the particular API call.

15. The API-authorizing system of claim 14, wherein: the set of servers distributes (i) a first storage structure storing a first set of policies and parameters to the first local API-authorizing agent and (ii) a second, different storage structure storing a second, different set of policies and parameters to the second local API-authorizing agent; and the first local API-authorizing agent requires a different set of API-authorization policies and parameters to authorize the API calls received by the first application than the second local API-authorizing requires to authorizes the API calls received by the second application.

16. The API-authorizing system of claim 14, wherein: the first application, the first plurality of additional applications, and the first local API-authorizing agent execute within a first virtual machine or container on the particular computer; and the second application, the second plurality of additional applications, and the second local API-authorizing agent execute within a second virtual machine or container on the particular computer.

17. The API-authorizing system of claim 14, wherein the server set distributes API-authorization policies and parameters to a plurality of computers executing machines for a particular tenant in a multi-tenant datacenter.

18. The API-authorizing system of claim 14, wherein the server set distributes API-authorization policies and parameters to a plurality of computers executing machines for a plurality of different tenants in a multi-tenant datacenter.
