

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12386714

Kind Code

B2

Date of Patent

August 12, 2025

Inventor(s)

Kothiyal; Rachita et al.

Resync transfer for recovering from storage site failure utilizing background pull operations

Abstract

Techniques are provided for performing a resync transfer to recover from a storage site failure. During normal operation of a first site hosting a first volume, data is replicated to a second volume hosted by a second site. If the first site fails, when clients are redirected to the second volume at the second site. When the first site recovers, data modifications made to the second volume are resynced back to the first volume. As part of synchronizing the first volume, a data warehouse is rebuilt at the first site in order to track the location of blocks present on the replication destination. Typically, the data modifications are transferred after the data warehouse is rebuilt, which results in significantly long resync times. The techniques provided herein decrease the resync time by either rebuilding the data warehouse in parallel with resyncing the data modifications or circumvent the need for rebuild.

Inventors: Kothiyal; Rachita (San Jose, CA), Pandit; Atul Ramesh (Los Gatos, CA), Naidu; Abhishek (San Jose, CA), Ponnapur; Anil Kumar (Sunnyvale, CA), George; Tijin (San Jose, CA)

Applicant: NetApp, Inc. (San Jose, CA)

Family ID: 96024144

Assignee: NetApp, Inc. (San Jose, CA)

Appl. No.: 18/541507

Filed: December 15, 2023

Prior Publication Data

Document Identifier

Publication Date

US 20250199923 A1

Jun. 19, 2025

Publication Classification

Int. Cl.: G06F11/16 (20060101)

U.S. Cl.:

CPC G06F11/1662 (20130101); G06F2201/815 (20130101)

Field of Classification Search

CPC: G06F (11/1662); G06F (2201/815)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7308545	12/2006	Kekre	711/161	G06F 11/2097
8010758	12/2010	Bezbaruah et al.	N/A	N/A
10360099	12/2018	Subramanian et al.	N/A	N/A
10521143	12/2018	Subramanian et al.	N/A	N/A
11036420	12/2020	Thoppil et al.	N/A	N/A
11048430	12/2020	Thoppil et al.	N/A	N/A
11132339	12/2020	Kaushik et al.	N/A	N/A
11151162	12/2020	Narasingarayanapeta et al.	N/A	N/A
11210013	12/2020	Thoppil et al.	N/A	N/A
11615001	12/2022	Naidu et al.	N/A	N/A
11650886	12/2022	Mathew et al.	N/A	N/A
11853104	12/2022	Naidu et al.	N/A	N/A
2014/0344526	12/2013	Brown	711/141	G06F 11/1458
2017/0161308	12/2016	Jarvis	N/A	G06F 16/2237
2021/0240348	12/2020	Xiong	N/A	G06F 3/065
2023/0029465	12/2022	Mallipeddi et al.	N/A	N/A
2023/0409535	12/2022	Kashi Visvanathan	N/A	G06F 21/602

OTHER PUBLICATIONS

“FabricPool tier management”, NetApp, URL: https://docs.netapp.com/us-en/ontap/pdfs/sidebar/FabricPool_tier_management.pdf, Apr. 16, 2024, 58 pages. cited by applicant

“HA pair management”, NetApp, URL: https://docs.netapp.com/us-en/ontap/pdfs/sidebar/HA_pair_management.pdf, Apr. 16, 2024, 26 pages. cited by applicant

“SVM data mobility”, NetApp, URL: https://docs.netapp.com/us-en/ontap/pdfs/sidebar/SVM_data_mobility.pdf, Apr. 16, 2024, 14 pages. cited by applicant

International Search Report and Written Opinion for Application No. PCT/US2024/060003, mailed on Apr. 22, 2025, 19 pages. cited by applicant

Background/Summary

BACKGROUND

(1) Many storage environments implement data replication and/or other redundancy data access techniques for data loss protection and non-disruptive client access. For example, a first storage site provides clients with primary access to data stored within a first volume. The data is replicated to a second volume maintained at a second storage site. If there is a failure at the first storage site, then clients are switched over to access the data from the second volume at the second storage site. This is performed as a switchover operation from the first storage site to the second storage site. If the first storage site recovers from the failure, then a switchback operation is performed. As part of the switchback operation, the first volume is resynced with the second volume by a resync operation because data modifications may have occurred at the second volume while the first volume was inaccessible. Once the first volume is resynchronized with the second volume, the clients are switched back to accessing the data through the first volume.

Description

DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1A is a block diagram illustrating an example of performing data replication from a first site to a second site in accordance with an embodiment of the present technology.
- (2) FIG. 1B is a block diagram illustrating an example of performing a failover operation in accordance with an embodiment of the present technology.
- (3) FIG. 1C is a block diagram illustrating an example of performing a resync operation in accordance with an embodiment of the present technology.
- (4) FIG. 2 is a flow chart illustrating an example method for performing a resync transfer to recover from a site failure in accordance with an embodiment of the present technology.
- (5) FIG. 3 is a block diagram illustrating an example of a system for performing a resync transfer to recover from a site failure in accordance with an embodiment of the present technology.
- (6) FIG. 4 is a flow chart illustrating an example method for performing a resync transfer to recover from a site failure in accordance with an embodiment of the present technology.
- (7) FIG. 5 is a block diagram illustrating an example of a system for performing a resync transfer to recover from a site failure in accordance with an embodiment of the present technology.
- (8) FIG. 6 is a flow chart illustrating an example method for performing a resync transfer to recover from a site failure in accordance with an embodiment of the present technology.
- (9) FIG. 7 is a block diagram illustrating an example of a system for performing a resync transfer to recover from a site failure in accordance with an embodiment of the present technology.
- (10) FIG. 8 is a block diagram illustrating an example of a node in accordance with an embodiment of the present technology.
- (11) FIG. 9 is an example of a computer readable medium in which an embodiment of the present technology may be implemented.

DETAILED DESCRIPTION

(12) Some examples of the claimed subject matter are now described with reference to the drawings, where like reference numerals are generally used to refer to like elements throughout. In the following description, for purposes of explanation, numerous specific details are set forth to provide an understanding of the claimed subject matter. It may be evident, however, that the

claimed subject matter may be practiced without these specific details. Nothing in this detailed description is admitted as prior art.

(13) A storage environment stores data across various storage sites such as for data redundancy, failover protection, and other types of data protection. The storage environment includes a first site (e.g., an on-premises computing environment, a cloud computing environment, a virtual machine, a storage virtual machine, a containerized environment, a storage controller, a node, or other hardware and/or software implemented as a primary storage site) that hosts a first volume as a primary volume that is made accessible for client access. A second site (a secondary storage site) hosts a second volume maintained as a secondary volume that is a backup copy of the primary volume. Data of the primary volume is replicated to the secondary volume so that the secondary volume mirrors the primary volume. Thus, if the first site fails, then clients are redirected to the secondary volume at the second site as part of a failover operation for non-disruptive access to client data.

(14) During a disaster where the first site has a failure and the first volume is not available to clients for accessing the client data, the second site performs a promote operation and transitions the second volume from being the secondary volume to being the primary volume now available for client access through the second site. In this way, the clients can now access the client data through the second volume that has been promoted as the primary volume as part of disaster recovery for the first site. Over time, data modifications are made to the second volume that are not replicated to the first volume because of the failure of first site.

(15) Once the first site recovers, a resync operation (a reverse resync transfer) is performed to replicate and protect the data modifications made to the second volume by transferring the data modifications to the first volume (currently demoted as the secondary volume). Eventually, client access is restored to first volume that is then promoted back to being the primary volume, and the second volume is demoted back to being the secondary volume that is a data protection copy of the primary volume. The resync operation is the process of bringing the primary and secondary volumes in sync where the primary and secondary volumes are storing the same data so that data protection can be restored for the first volume because the first volume and the second volume mirror one another after the resync operation.

(16) A replication engine that replicates data from the primary volume to the secondary volume maintains a data warehouse for the secondary volume. The data warehouse is a database or map of data blocks present on the secondary volume. That is, the data warehouse tracks the locations of each data block on the secondary volume (e.g., tracking virtual volume block numbers such as a virtual volume block number of a data block on the secondary volume that is mapped to a virtual volume block number of a corresponding block on the primary volume that stores the same data). The data warehouse is used to deduplicate and share data blocks because the data warehouse can be used to track when there are multiple references to the same data block (e.g., multiple files referring to the same data block).

(17) During the disaster recovery, the replication engine processes the data warehouse at the second site to create virtual volume block number (VVBN) mappings from the second volume (primary volume) to the first volume (secondary volume). The mappings are transmitted as replication operations from to the first site, and are populated as entries within the data warehouse being rebuilt. This enables the ability to share data blocks (e.g., sharing between a data warehouse file block number and a container file of the first volume) such that further transfers can share the existing blocks in order to preserve storage efficiency (e.g. if a data block of data is being resynced to the first volume that already comprises that data, then the data is not redundantly stored again but a reference is created to indicate that the data is now pointed to by multiple sources).

(18) Unfortunately, data sharing consumes a significant amount of time during the resync operation. The replication engine performs resync transfers that would have to rebuild the entire data warehouse at the first site for the first volume (currently demoted to being the secondary

volume) that will be updated with the data modifications of the second volume that is currently promoted to being the primary volume. Because the entire data warehouse must be rebuilt for the entire volume and not just for the data modifications, the resync transfer times are proportional to the size of the entire volume. Because volumes can be tends to hundreds of TBs, it is difficult if not impossible for the replication engine to satisfy disaster recovery service level objectives. Thus, the rebuilding of the data warehouse is a long running process that can take a substantial amount of time (e.g., hours to days) that exceeds the disaster recovery service level objects of clients. Furthermore, the data warehouse is rebuilt before the data modifications are transferred from the second volume to the first volume, which also increases the overall resync operation. Because the replication engine can be used as a common transfer engine for other processes such as asynchronous replication, synchronous replication, and Vserver migration use cases, the longer resync times become a broad problem.

(19) FIGS. 1A-1C are block diagrams illustrating an example of a system **100** for performing data replication from a first site **102** to a second site **104**, a failover operation, and resync and data warehouse rebuilt operations. The first site **102** hosts a first volume **106** as a primary volume **118** that is made accessible to clients such as for processing I/O operations **112** from a client **110**, as illustrated by FIG. 1A. The second site **104** may maintain a second volume **108** as a secondary volume **120** that is a standby backup for the first volume **106**. In this way, data replication **114** is performed to replicate data from the first volume **106** to the second volume **108** (e.g., synchronous replication, asynchronous replication, etc.). A data warehouse **116** is maintained at the second site **104** for the secondary volume **120**. The data warehouse **116** is used to track the location of blocks of the secondary volume **120**. The data warehouse **116** may include mappings for the blocks of the secondary volume **120** such as a mapping of virtual volume block numbers used by the primary volume **118** to virtual volume block numbers used by the secondary volume **120**.

(20) The first site **102** may experience a failure **121** such that the first volume **106** is no longer available to the clients, as illustrated by FIG. 1B. Accordingly, a failover operation **122** is performed to failover the clients such as I/O operations **124** of the client **110** to the second volume **108** hosted at the second site **104**. As part of the failover operation **122**, the second volume **108** is promoted to be the primary volume **118** and the first volume **106** will be demoted to be the secondary volume **120**. In this way, clients are provided with non-disrupted access to client data through the second site **104** and the second volume **108**.

(21) The first site **102** recovers from the failure, as illustrated by FIG. 1C. The second volume **108** is the primary volume **118** and the first volume **106** is the secondary volume **120**. In order to eventually switch the clients back to accessing client data through the first volume hosted by the first site **102**, a resync operation **130** will be performed to synchronize data modifications (e.g., the second volume **108** may be modified while the second volume **108** is the primary volume **118** being accessed by the clients) of the second volume **108** to the first volume **106** in order to synchronize the first volume **106** to the second volume **108**. As part of resyncing the first volume **106** with the data modifications of the second volume **108**, a data warehouse **134** is rebuilt, by a data warehouse rebuild operation **132**, at the first site **102** for tracking mappings of virtual volume block numbers between the first volume **106** and the second volume **108** so that block sharing and deduplication is implemented for the first volume **106**. Various embodiments of the overall resync transfer from the second volume **108** to the first volume **106** are further discussed in relation to FIGS. 2-7.

(22) The techniques provide herein are capable of reducing the time to perform the resync transfer from hours/days down to minutes. Instead of rebuilding/populating the data warehouse at the first site and waiting to transfer the data modifications from the second volume to the first volume until the data warehouse is rebuilt, the disclosed technique logs virtual volume block number mappings into a log metafile in parallel with the data modifications being resync from the second volume to the first volume.

(23) In some embodiments of the disclosed techniques, the resync transfer is initiated. During the resync transfer, virtual volume block number mappings (e.g., mappings from virtual volume block numbers of the second volume to virtual volume block numbers of the first volume) are logged into a log metafile. In some embodiments, the log metafile is a flat file that is indexed by virtual volume block numbers of the second volume, with a payload as virtual volume block numbers from the first volume. The indexes are at fixed offsets, which provides for the ability to identify an index very quickly such as within a constant time. In some embodiments, a maximum number of indices is set or dependent on a largest virtual volume block number from the second volume. In some embodiments, the log metafile is a volume metafile that is captured by snapshots of a volume such as the first or second volume. Once all virtual volume block number mappings are log into the log metafile, the resync transfer is declared complete. At this point, transfers of the data modifications from the second volume to the first volume can proceed. This reduces the delay in starting the data transfers because the entire data warehouse does not need to be rebuilt before the data modification transfers can occur.

(24) A scanner (e.g., a background scanner) is implemented to process (drain) virtual volume block number mappings (entries) from the log metafile. The scanner processes the entries within the log metafile asynchronously with respect to the resync transfer that is transferring blocks of data modifications from the second volume to the first volume. For each mapping from a virtual volume block number of the second volume to a virtual volume block number of the first volume, a block share is performed such as between a data warehouse file block number of the second volume (a data warehouse FBN) and a container file block number of the first volume. A container file holds blocks of a volume, and thus the container file block number references the container file holding the blocks of the first volume. This block sharing is essentially populating entries in the data warehouse in order to rebuild the data warehouse at the first site. Once complete, the log metafile is deleted.

(25) Various operations are implemented for the log metafile. In some embodiments, a look up to the log metafile is performed by a push replication operation (a replication operation). When the data modification transfers start/resume from the second volume to the first volume and replication operations are received, a lookup is performed within the data warehouse and also the log metafile. If an entry is located, then block sharing with that corresponding container file block number is performed. The entry within the log metafile can be retained because the scanner will subsequently process the entry while rebuilding the data warehouse. Thus, there is merely a constant time lookup cost to the log metafile for the resync transfers. In some embodiments, file buffers are loaded into memory on an as needed basis. Given that the log metafile is relatively dense with mappings, locality is leveraged for cache/memory.

(26) In some embodiments, a push name free operation is performed for when a virtual volume block number is no longer is use. For example, a resync transfer (a rebuild transfer) from the second volume to the first volume is performed using a first snapshot (S1). In this example, the snapshot (S1) has a first file referring to a first virtual volume block number (V1). The log metafile has been populated, but not yet processed (drained) where mappings (entries) within the log metafile have not yet been fully processed to rebuild the data warehouse. A second snapshot (S2) captures a state where the first file has been deleted. As such, the virtual volume block number (V1) is freed within the second snapshot (S2). The second snapshot (S2) is transferred from the second volume to the first volume. At this point, the log metafile is still not fully processed (drained). A push name free operation is received for the virtual volume block number (V1), and a push inode operation to delete the first file and remove references of the virtual volume block number (V1) is received. First, the data warehouse on the first site is checked. If a mapping is identified, then the push name free operation is processed to free the virtual volume block number (V1). Next, the log metafile is checked. If the log metafile has a mapping for the virtual volume block number (V1), then the mapping is invalidated by writing a virtual block number (VBN) hole

as a payload.

(27) In some embodiments, a push name data operation is performed. In an example, a virtual volume block number of the push name data operation is a new allocation of a block in an incremental resync transfer and was not previously allocated. Accordingly, there is no existing mapping in the data warehouse nor in the log metafile. Logic for processing the push name data operation is configured to handle this situation. In another example, a virtual volume block number is freed and reallocated. For example, a rebuild (resync) transfer from the second site to the first site is performed using a first snapshot (S1). In this example, the first snapshot (S1) has a first file and a second file that both refer to a virtual volume block number (V1). The log metafile is populated with mappings, but has not yet been fully processed (drained). A second snapshot (S2) captures a state where the first file has been deleted. So, only the second file is referring to the virtual volume block number (V1) on the second volume. The second snapshot (S2) is transferred from the second site to the first volume. At this point, the log metafile is still not fully processed (drained). Now, the second file is deleted and the first snapshot (S1) is deleted on the second site. The virtual volume block number (V1) is now free on the second volume. A new file (a third file) is created, and the virtual volume block number (V1) is allocated again for use by the new file. A third snapshot (S3) is created and transferred to the first site. This will result in a push name data operation for the virtual volume block number (V1). There could be an entry for the virtual volume block number (V1) in the data warehouse or in the log metafile, or both. An update is performed for the data pointed to by the virtual volume block number (V1) without allocating a new virtual volume block number. A new virtual volume block number is not allocated so that any replication operations that also arrived in the transfer at any point in time can use a valid mapping. This is because if a mapping exists in both the data warehouse and the log metafile, the mappings are to match.

(28) In some embodiments, a push name reuse operation is performed. When the push name reuse operation is received, a block share is performed with a virtual volume block number from an older snapshot. The data warehouse is evaluated from the older snapshot in order to resolve a mapping. If the mapping does not exist, then the log metafile is checked to resolve the mapping.

(29) A snapshot difference operation may interact with the log metafile. If a new resync (a flip resync now from the first site and first volume to the second site and the second volume), then the snapshot difference operation generates mappings (references) from the first volume to the second volume, and sends the mappings to the second site. In addition to evaluating the data warehouse for the mappings, the snapshot difference operation evaluates the log metafile, if the log metafile exists, and also transfers the mappings within the log metafile.

(30) In some embodiments, transfers are started without waiting for the rebuild of the data warehouse. The rebuild transfer to rebuild the data warehouse could be executed as a separate scanner in parallel with the data (data modification transfer) and reference (mappings transfer) phase of the overall resync transfer. This allows the transfer to start immediately on the flip resync and there is no delay for the data warehouse rebuild phase to complete.

(31) Data warehouse mappings could be logged in parallel with the processing of other operations for the overall resync transfer. A rebuild scanner that rebuilds the data warehouse operates on a reference snapshot, while the transfer pertains to incremental changes (data modifications) between the reference snapshot and a transfer snapshot. The rebuild scanner is operating in the background to process (drain) the mappings (entries) from the log metafile to create mappings (entries) in the data warehouse. In some embodiments, the transfer is declared complete when the rebuild scanner finishes logging the mappings to the log metafile, and the data (data modification transfer) and reference (mappings transfer) phase are complete. In some embodiments, the rebuild scanner includes both logging and reconciliation functionality that occur in parallel with the data modification transfers such that a current or future data modification data transfer does not wait on the rebuild scanner to either start or finish. In this way, the time taken for end-to-end transfer

completion merely includes the data and reference phase, and remains completely independent of the resync rebuild phase of the data warehouse (e.g., part of the transfer can vary in transfer time depending upon an amount of logging rebuild that is done based upon data set sizes). This provides a technical solution for making a resync transfer time dependent only on incremental snapshot delta (e.g., data modifications as opposed to data of the entire volume).

(32) The rebuild scanner executes in parallel with transfers that occur only for the first transfer done as part of the flip resync transfer. This parallel scenario may affect replay operations that interact with the log metafile during the first transfer. In some embodiments, a push replication operation (replay operation) is implemented. A replication operation first checks the data warehouse for a mapping. If the mapping does not exist, then the log metafile is checked for the mapping. If the mapping is also not found in the log metafile such as because the rebuild transfer has not completed, then the replication operation is buffered (e.g., queued into a queue). Once the data phase (data modification transfer/resync) for the transfer finishes and the rebuild phase to rebuild the data warehouse finished, then the buffered replication operation and/or other buffered replication operations are replayed. As part of replay, the data warehouse is first checked, and then the log file is checked. If a match is found, then block sharing is implemented. Otherwise, the block does not already exist at the first site and is pulled from the second site.

(33) In some embodiments, a push name free operation may be implemented. If the push name free operation arrives before the rebuild transfer inserts a mapping into the data warehouse, then it cannot be determined as to which local virtual volume block number to free. Accordingly, a VBN hole is inserted as a mapping into the log metafile. Subsequently when the mapping is obtained as part of the rebuild scanner, the VBN hole overrides the incoming mapping. If there is a mapping in the log metafile when the push name free operation is received, then the push name free operation is processed and the mapping is invalidated by inserting a VBN hole in place of the mapping. In some embodiments, a push name data operation is received for the first transfer, and belongs to new blocks that are allocated between the reference snapshot (a common snapshot capture a state of the first volume and the second volume when both values are storing the same data) and the incremental snapshot of the secondary volume. Because there is no dependency between the push name data operation and the log metafile, the push name data operation is processed such that a block is allocated and shared with the data warehouse (a data warehouse file used to store/represent the data warehouse).

(34) FIG. 2 is a flow chart illustrating an example method **200** for performing a resync transfer to recover from a site failure, which is described in conjunction with system **300** of FIG. 3. During operation **202** of method **200**, data is replicated from the first volume **106** hosted at the first site **102** as a primary volume to the second volume **108** hosted at the second site **104** as a secondary volume. While the first volume **106** is hosted as the primary volume, client I/O operations are directed to the primary volume. In some embodiments, the client I/O operations or data of the first volume **106** is replicated (e.g., synchronous replication or asynchronous replication) to the second volume **108** so that the second volume **108** mirrors the first volume **106**. The data warehouse **116** is maintained at the second site **104** for the secondary volume **120**. The data warehouse **116** is used to track the location of blocks of the secondary volume **120**. The data warehouse **116** is populated with mappings of virtual volume block numbers used by the primary volume such as the first volume **106** to virtual volume block numbers used by the secondary volume such as the second volume **108**.

(35) A failure may occur at the first site **102** such that the first volume **106** is no longer accessible for client access. In response to the first site **102** experiencing the failure, the second volume **108** is promoted to be the primary volume and client I/O operations are redirected to the second volume **108**, during operation **204** of method **200**. Some of the client I/O operations (e.g., write operations, delete operations, create operations, etc.) may result in data modifications to the second volume **108** such that the second volume **108** diverges from the first volume **106**. In this way, the data

warehouse **116** tracks virtual volume block numbers of blocks storing data of the second volume **108** in relation to virtual volume block numbers of corresponding blocks of the first volume **106**. (36) At a subsequent point in time, the first site **102** recovers such that the first volume **106** becomes available for resynchronization with the second volume **108**, as illustrated by FIG. 3. In response to the first site **102** recovering, a resync operation **302** is initiated to transfer the data modifications from the second volume **108** to the first volume **106** so that the first volume **106** can be re-promoted to be the primary volume after being resynced, during operation **206** of method **200**. During the resync operation **302**, database rebuild operations are logged into a log metafile **308**, during operation **208** of method **200**. The database rebuild operations are used to rebuild a database file **306** used to track blocks that are shared by files of the first volume **106** and/or the second volume **108**. In some embodiments, virtual volume block numbers of the second volume **108** are received from the second site **104** as part of the database rebuild operations. The virtual volume block numbers are logged into the log metafile **308** that is indexed by the virtual volumes block numbers. The indexes are at fixed offsets and a maximum number of indices is derived from a largest virtual volume block number from the second site **104** for the second volume **108**. In some embodiments, virtual volume block numbers from the second site **104** are logged into the log metafile indexed by the virtual volume block numbers and are mapped to virtual volume block numbers used by the first site **102**. In this way, virtual volume block number mappings between the first volume **106** and the second volume **108** are logged within the log metafile **308**. The log metafile **308** can be evaluated to determine whether a block has been processed as part of the resync operation **302** (e.g., whether a block comprising a data modification has been replicated from the second volume **108** to the first volume **106**).

(37) During operation **210** of method **200**, a background scanner **310** is implemented to reconcile information within the log metafile **308** for rebuilding the database file **306**. The background scanner **310** is implemented to process (drain) virtual volume block number mappings (entries or map entries) from the log metafile **308**. The background scanner **310** processes the entries within the log metafile **308** asynchronously such as with respect to the resync operation **302**. For each mapping from a virtual volume block number of the second volume **108** to a virtual volume block number of the first volume **106**, a block share is performed such as between a data warehouse file block number of the second volume **108** (a data warehouse FBN) and a container file block number of the first volume **106**. This block sharing is essentially populating entries in the database file **306** in order to rebuild a data warehouse **304** at the first site **102**.

(38) In some embodiments, the data warehouse **304** is represented by both the database file **306** and the log metafile **308** until the background scanner **310** has finished. In this way, both the log metafile **308** and the database file **306** are used together to obtain a consistent state of the first volume **106** and/or the second volume **108** (e.g., used to identify deduplicated data such as where multiple files share the same block of data). For example, the log metafile **308** and the database file **306** are used to deduplicate blocks storing data of the first volume **106** such as by tracking what blocks are referenced by more than one file. In some embodiments, the background scanner **310** asynchronously processes mappings (entries or map entries) within the log metafile **308**. For a map entry that maps a virtual volume block number from the second volume **108** to a virtual volume block number of the first volume **106**, a block share is performed between a file block number (e.g., a file block number of the data warehouse **116** of the second volume **108**) to a container file block number associated with the first volume **106**.

(39) Once the background scanner **310** has completed, the log metafile **308** is deleted, and the data warehouse **304** is now represented by the database file **306** (e.g., the database file **306** becomes the data warehouse **304**). The data warehouse **304** is used to track locations of blocks within the first volume **106** and/or the second volume **108** such as by mapping virtual volume block numbers used by the primary volume such as the second volume **108** to virtual volume block numbers used by the secondary volume such as the first volume **106**.

(40) In response to the resync operation **302** completing, the second volume **108** is demoted to become the secondary volume and the first volume **106** is promoted to become the primary volume for client access. The resync operation **302** may be specified as complete in response to all virtual volume block numbers from the second site **104** being logged into the log metafile **308**, thus all virtual volume block numbers are accounted for between the log metafile **308** and the database file **306** as being reconciled by the background scanner **310**.

(41) In some embodiments of the disclosed techniques, virtual volume block number mappings are embedded within replication operations during transfer instead of performing a separate data warehouse rebuild phase. That is, instead of implementing a separate data warehouse rebuild phase for sending over mappings of virtual volume block numbers between the secondary volume and the primary volume, such information is embedded into operations/commands (e.g., replication operations, replay operations, etc.) during the transfer.

(42) In an example where the replication relationship is flipped, transfers are being performed from the second volume to the first volume. A data warehouse is located at the second site, and is captured in a reference snapshot that has mappings of virtual volume block numbers from the first volume to virtual volume block numbers of the second volume. During the transfer from the second volume to the first volume, replication operations are generated to describe an inode number, an offset, and virtual volume block numbers of the second volume. As provided herein, corresponding virtual volume block numbers of the first volume can be also be included within the replication operations to map to the virtual volume block numbers of the second volume. This provides the ability for block sharing with the virtual volume block numbers of the first volume when the replication operation arrives, without the need to have a data warehouse at the second site (the destination site of the replication operation). However, the existing mappings at the second site are from the first volume to the second volume. Accordingly, an efficient lookup of mappings from the second volume to the first volume is performed. Some virtual volume block numbers of the second volume may not have mappings within the data warehouse. These virtual volume block numbers are not present on the first volume, and were newly allocated in the current incremental transfer. Accordingly, a push name data operation is sent for these virtual volume block numbers for block sharing.

(43) FIG. 4 is a flow chart illustrating an example method **400** for performing a resync transfer to recover from a site failure, which is described in conjunction with system **500** of FIG. 5. During operation **402** of method **400**, data is replicated from the first volume **106** hosted at the first site **102** as a primary volume to the second volume **108** hosted at the second site **104** as a secondary volume. While the first volume **106** is hosted as the primary volume, client I/O operations may be directed to the primary volume. In some embodiments, the client I/O operations or data of the first volume **106** is replicated (e.g., synchronous replication or asynchronous replication) to the second volume **108** so that the second volume **108** mirrors the first volume **106**. The data warehouse **116** is maintained at the second site **104** for the secondary volume **120**. The data warehouse **116** is used to track the location of blocks of the secondary volume **120**. The data warehouse **116** is populated with mappings of virtual volume block numbers used by the primary volume such as the first volume **106** to virtual volume block numbers used by the secondary volume such as the second volume **108**.

(44) A failure may occur at the first site **102** such that the first volume **106** is no longer accessible for client access. In response to the first site **102** experiencing the failure, the second volume **108** is promoted to be the primary volume and client I/O operations are redirected to the second volume **108**, during operation **404** of method **400**. Some of the client I/O operations (e.g., write operations, delete operations, create operations, etc.) may result in data modifications to the second volume **108** such that the second volume **108** diverges from the first volume **106**. In this way, the data warehouse **116** tracks virtual volume block numbers of blocks storing data of the second volume **108** in relation to virtual volume block numbers of corresponding blocks of the first volume **106**.

(45) At a subsequent point in time, the first site **102** recovers such that the first volume **106** becomes available for resynchronization with the second volume **108**, as illustrated by FIG. 5. In response to the first site **102** recovering, a resync operation **502** is initiated to transfer the data modifications from the second volume **108** to the first volume **106** so that the first volume **106** can be re-promoted to be the primary volume after being resynced.

(46) During the resync operation **502**, the data warehouse **116** at the second site **104** is evaluated to identify mappings between virtual volume block numbers, used by the first volume **106** to reference blocks storing data, to virtual volume block numbers used by the second volume to reference blocks storing corresponding data (corresponding data that is a replica of the data of the first volume **106**), during operation **406** of method **400**.

(47) During operation **408** of method **400**, reverse mapping generation **504** is implemented to generate reverse mappings. The reverse mappings may map the virtual volume block identifiers used by the second volume **108** to the virtual volume block identifiers used by the first volume **106**. In some embodiments, the reverse mapping generation **504** generates the reverse mappings using the mappings from the data warehouse **116**.

(48) In some embodiments, the data warehouse **116** at the second site **104** is a flat file that is indexed by virtual volume block numbers of the first volume **106**. Thus, mappings (entries) within the data warehouse **116** at the second site **104** map the virtual volume block numbers of the first volume **106** to virtual volume block numbers of the second volume **108** that store corresponding data (e.g., a virtual volume block number of the first volume **106** refers to particular data, and is mapped to a virtual volume block number of the second volume **108** that refers to a replicated copy of that data). At the start of the resync operation **502**, the data warehouse **116** (e.g., the flat file that is indexed by virtual volume block numbers of the first volume **106**) is traversed (walked) to start building a reverse index of reverse mappings. In some embodiments, the data warehouse **116** is read in a linear fashion to build the reverse index as a separate metafile to track reverse mappings. The reverse index may be trained until the resync completes. The time to build the reverse index (metafile) is $O(n)$ where n is a total number of mappings (reverse mappings). Since entries (reverse mappings) are only a few bytes each (e.g., 8 bytes) and are condensed, the $O(n)$ is very quick and efficient. The reverse index of mappings map the virtual volume block numbers of the second volume **108** to virtual volume block numbers of the first volume **106** that store corresponding data. In some embodiments, the reverse index is implemented as a persistent key/value pair mapping file. In some embodiments, the reverse index is efficiently built because the data warehouse **116** is only traversed once. Additionally, lookups of mappings within the reverse index can be performed at a near constant time since merely a lookup of a virtual volume block number is performed at a given index.

(49) Once the reverse mappings are created (e.g., created within the reverse index such as the persistent key/value pair mapping file), the transfer by the resync operation **502** of data modifications from the second volume **108** to the first volume **106** will replicate any changes from the second volume **108** (e.g., the data modifications) to the first volume **106**. As part of the transfer, new data blocks of the data modifications are transmitted to the first site **102** to apply to the first volume **106** during a data phase. The new data blocks are transferred and written through a data warehouse file (e.g., a data warehouse at the first site **102** for the first volume **106**) during the data phase. The references to these new data blocks and old data blocks (e.g., virtual volume block numbers of the new data blocks and the old data blocks) are transferred in a reference phase. As part of the reference phase, a replication operation is sent from the second site **104** to the first site **102** for each reverse mapping. The replication operation for a data block is updated to include a corresponding virtual volume block number of the first volume **106** used to reference the data block at the first site **102** (e.g., populated with or based upon a reverse mapping). The operation may specify blocks used to store data of a file so that the first site **102** can track and map to the blocks used by the second site **104** to store the data of the file. The corresponding virtual volume

block number of the first volume **106** used to reference the data block is obtained from a reverse mapping within the reverse index. The corresponding virtual volume block number within the replication operation enables the first site **102** to perform block sharing at the first volume **106** without having to rebuild the data warehouse for the entire data of the first volume **106**.

(50) As an example, a snapshot (S1) of the first volume **106** can include a first data block with a virtual volume block number (1), a second data block with a virtual volume block number (2) and a third data block with a virtual volume block number (3), and there is 1 file referencing these data blocks. The data blocks are transferred from the first site **102** to the second site **104**, which are stored at data blocks having a virtual volume block number (10), a virtual volume block number (20), and a virtual volume block number (30). When the first site **102** fails, clients are provided with access to the second volume **108** at the second site **104**. A new data block with a virtual volume block number (9) is allocated at the second volume **108**, and a new file is created using the virtual volume block number (10), the virtual volume block number (20), and the virtual volume block number (30), which still exist at the first volume **106** (e.g., as virtual volume block numbers (1)-(3)). During a resync from the second site **104**, the virtual volume block number (9) is transferred to the first site **102** and a mapping is created within a data warehouse that maps virtual volume block number (9) to a new block for the first volume **106**, during a data phase. During a reference phase, a replication operation is sent with a reverse mapping. In this way, the replication operation includes a first entry that comprises a name of the new file, a file block number (1), and a mapping between the virtual volume block number (10) and the virtual volume block number (1); a second entry that comprises the name of the new file, a file block number (2), and a mapping between the virtual volume block number (20) and the virtual volume block number (2); and a third entry that comprises the name of the new file, a file block number (3), and a mapping between the virtual volume block number (30) and the virtual volume block number (3). The virtual volume block numbers (1)-(3) are used to perform block sharing with existing data of the first volume **106** without having to rebuild the data warehouse for the entire dataset.

(51) During operation **410** of method **400**, a command **506** is generated for each block being sent from the second site **104** to the first site **102** (e.g., a block of data modifications being resynced from the second volume **108** to the first volume **106** by the resync operation **502**). The command **506** comprises the data of the block being resynced from the second volume **108** to the first volume **106** and a reverse mapping for the block. In some embodiments, the command **506** may be constructed as a replication operation that describes an inode (inode number of a file), an offset, a virtual volume block identifier of the second volume **108**, and/or other information. In some embodiments, the command **506** may be constructed as a replication operation that describes an inode (inode number of a file), an offset, virtual volume block numbers of the second volume **108**, and corresponding virtual volume block numbers of the first volume **106**. In some embodiments, the replication operation is constructed during the resync operation **502**.

(52) In some embodiments, the reverse mapping generation **504** may determine that the data warehouse **116** lacks a mapping for a virtual volume block number used by the second volume **108**. In an example, the data warehouse **116** may lack the mapping for the virtual volume block number because the virtual volume block number is not currently being used at the first site **102**. In another example, the data warehouse **116** may lack the mapping for the virtual volume block number because the virtual volume block number was newly allocated as part of an incremental transfer from the second site **104** to the first site **102**. Accordingly, a push name data command is transmitted to the first site **102** for performing a share operation.

(53) The command **506** is transmitted to the first site **102** for updating the first volume **106**. The data of the block may be stored for the first volume **106**. A block sharing operation **508** may be performed to share the block based upon the reverse mapping embedded within the command **506** so that a single instance of the data is stored and can be referenced by any number of times by any number of files (e.g., a reference count for a block of data may be incremented each time the data is

referenced, and is decremented each time the data is no longer reference).

(54) Reverse mappings within the commands transmitted from the second site **104** to the first site **102** during the resync operation **502** are used to perform the block sharing operation **508** in order to deduplicate the first volume **106** (e.g., multiple files can point to a same block of data that is stored once instead of multiple times).

(55) In response to the resync operation **502** completing, the second volume **108** is demoted to become the secondary volume and the first volume **106** is promoted to become the primary volume.

(56) In some embodiments of the disclosed techniques, a background pull based data warehouse rebuild is implemented. The background pull based data warehouse rebuild is similar to performing a quick resync of data modifications, and then the data warehouse is built in the background. A data warehouse scanner is implemented at the destination (e.g., a background scanner and/or a snapshot difference scanner). The data warehouse scanner (background scanner or snapshot difference scanner) does a logical scan of a common snapshot (e.g., a snapshot capturing the first volume and the second volume when both volumes store the same data) and performs the follow actions. The data warehouse scanner performs a snapshot difference operation for a common snapshot's public inofile (e.g., a common snapshot being a snapshot that is common between the first volume and the second volume, and the public inofile being an inofile listing inodes of files within the volumes) with a null snapshot. The data warehouse scanner issues a baseline snapshot difference operation (e.g., a compare with null operation) for each changed inode (e.g., an inode of a file that has changed at the second volume after the failure of the first site, thus resulting in a data modification), and a changed vector of information is determined. The changed vector of information includes an inode, a file block number, a destination virtual volume block number, a consistency group size (e.g., a grouping of data that is to be replicated together as a consistency group), etc. The data warehouse scanner issues a baseline snapshot difference operation (e.g., a compare with null operation) for certain files such as private files (special private files).

(57) As part of rebuilding the data warehouse (a data warehouse map), certain information in the entire lifecycle of the rebuild is tracked for each changed virtual volume block number reported by the snapshot difference scanner (e.g., a change virtual volume block number for a block of data of the second volume **108** modified since the failure of the first site **102**). Such information may include an inode, a file block number, a destination virtual volume block number, a consistency group size (e.g., a grouping of data that is to be replicated together as a consistency group), etc. For each entry of information (for each entry that includes an inode, a file block number, a destination virtual volume block number, a consistency group size, etc.), a source virtual volume block number is obtained from the source of the transfer such as the second site/second volume. Once a mapping (a virtual volume block number mapping) is received from the source (e.g., received from the second site), a determination is made as to whether the mapping is already present. If the mapping is already present, then no further action is taken. If a one-to-one mapping between source virtual volume block numbers and destination virtual volume block numbers cannot be maintained (e.g., due to differences in consistency group geometry or compression cases), an entry is not populated within the data warehouse. If the mapping is not already present, then an entry within the data warehouse is marked with the mapping.

(58) In some embodiments, a source snapshot of the source volume used for the transfer is locked. Thus, send blocks for a virtual volume block number cannot be received in a data transfer (e.g., a data transfer of a block storing data modifications of the second volume being resynced to the first volume) happening in parallel. Thus, a verification (or an indication is created) for this scenario in a send blocks code module to indicate that an entry within the data warehouse must be a hole while writing to the data warehouse via a send blocks operation.

(59) In some embodiments, race condition occurs with holes where a data warehouse rebuild virtual volume block number mapping of a virtual volume block number (e.g., rebuilding the data

warehouse with virtual volume block number mappings) races with the same virtual volume block number being transitioned to a hole as part of a parallel transfer. If the data warehouse mapping is applied first and then the hole is applied, then a consistent result is achieved. Otherwise, a new ondisk flag is set in an indirect of the data warehouse (e.g., an indirect block/point) for each entry while punching (removing/deleting) a hole as part of the transfer. The mapping is applied only if the flag is not set. If the data warehouse indirect is not allocated while punching the hole from the transfer, then an indirect is allocated and the flag is set within the indirect.

(60) FIG. 6 is a flow chart illustrating an example method **400** for performing a resync transfer to recover from a site failure, which is described in conjunction with system **700** of FIG. 7. During operation **602** of method **600**, data is replicated from the first volume **106** hosted at the first site **102** as a primary volume to the second volume **108** hosted at the second site **104** as a secondary volume. While the first volume **106** is hosted as the primary volume, client I/O operations may be directed to the primary volume. In some embodiments, the client I/O operations or data of the first volume **106** is replicated (e.g., synchronous replication or asynchronous replication) to the second volume **108** so that the second volume **108** mirrors the first volume **106**. The data warehouse **116** is maintained at the second site **104** for the secondary volume **120**. The data warehouse **116** is used to track the location of blocks of the secondary volume **120**. The data warehouse **116** is populated with mapping of virtual volume block numbers used by the primary volume such as the first volume **106** to virtual volume block numbers used by the secondary volume such as the second volume **108**.

(61) A failure may occur at the first site **102** such that the first volume **106** is no longer accessible for client access. In response to the first site **102** experiencing the failure, the second volume **108** is promoted to be the primary volume and client I/O operations are redirected to the second volume **108**, during operation **604** of method **600**. Some of the client I/O operations (e.g., write operations, delete operations, create operations, etc.) may result in data modifications to the second volume **108** such that the second volume **108** diverges from the first volume **106**. In this way, the data warehouse **116** tracks virtual volume block numbers of blocks storing data of the second volume **108** in relation to virtual volume block numbers of corresponding blocks of the first volume **106**.

(62) At a subsequent point in time, the first site **102** recovers such that the first volume **106** becomes available for resynchronization with the second volume **108**, as illustrated by FIG. 7. In response to the first site **102** recovering, a resync operation **702** is initiated to transfer the data modifications from the second volume **108** to the first volume **106** so that the first volume **106** can be re-promoted to be the primary volume after being resynced. During operation **606** of method **600**, a background scanner **706** is executed for rebuilding a data warehouse **708** at the first site **102**. The background scanner **706** is executed to track a lifecycle of rebuilding the data warehouse **708** by storing a changed vector of information **712**. The changed vector of information **712** is populated with entries, where an entry includes an inode of a file, a file block number, a destination virtual volume block number, and size information such as a consistency group size of a consistency group that includes data, of the file, stored at the destination virtual volume block number and referenced by the file block number.

(63) The background scanner **706** processes changed virtual volume block numbers of the second volume **108** identified by a snapshot difference scanner **704** using one or more snapshots **703**. The snapshot difference scanner **704** may compare a first snapshot (e.g., a common snapshot when both the first volume **106** and the second volume **108** were in-sync, and thus the common snapshot captures data common to both volumes) and a second snapshot (e.g., a current snapshot of the second volume **108** after data modifications have occurred) to identify the changed virtual volume block numbers corresponding to the data modifications made to the second volume **108** that are not reflected in the first volume **106**.

(64) In some embodiments, the snapshot difference scanner **704** is executed to generate a baseline difference by determining a difference between a null snapshot (e.g., a snapshot of a structure of a

volume but without data of the volume) and the common snapshot capturing data common between the first volume **106** and the second volume **108**. The snapshot difference scanner **704** may determine the difference as a difference between the null snapshot and a public inofile of the common snapshot (e.g., an inofile populated with inodes of files within a volume whose state is captured by the common snapshot). A baseline difference is issued for each changed inode identified for the second volume **108** (e.g., a changed inode being an inode of a file that changed based upon the data modifications to the second volume **108**, which is identified based upon the difference identified by the snapshot difference scanner **704**). In some embodiments, a baseline difference is issued for each changed inode identified for the second volume **108** to determine the changed vector of information **712** populated with the entries that each include an inode (e.g., an inode number of a changed file), a file block number, a destination virtual volume block number, and size information.

(65) For a changed virtual volume block number of the second volume **108**, the background scanner **706** obtains a source volume block number of the first volume **106** from the first site **102**, during operation **608** of method **600**. The source volume block number may be retrieved for an entry within the changed vector of information **712** for determining whether a new mapping is to be created as a new entry within the data warehouse **708**. The new mapping (the new entry) is created based upon whether the data warehouse **708** already comprises a mapping, within existing mappings **710**, between the changed virtual volume block number and the source virtual volume block number.

(66) During operation **610** of method **600**, the new entry is created as part of the background scanner **706** creating new entries **714** (new mappings) within the data warehouse **708** based upon the data warehouse **708** not already comprising a mapping between the changed virtual volume block number and the source virtual volume block number. The new entry is created to generate a mapping between the changed virtual volume block number and the source virtual volume block number. In this way, when the block of modified data is resynced by the resync operation **702** from the second volume **108** to the first volume **106**, the mapping is used to perform block sharing for any files referencing the modified data of the block. If the data warehouse **708** already comprises the mapping between the changed virtual volume block number and the source virtual volume block number, then the mapping is retained within the data warehouse **708** and no new entry/mapping is created.

(67) In some embodiments, a one to one mapping between source volume block numbers and destination volume block numbers cannot be maintained such as due to compression differences (e.g., data may be compressed into blocks that are arranged differently at the first site **102** and the second site **104**), consistency group geometry differences (e.g., a structure of how files and directories of a consistency group at the first site **102** may be different than how the files and directories are structured at the second site **104**). If the one to one mapping between the source volume block numbers and the destination volume block numbers cannot be maintained, then the new entry is not populated within the data warehouse **708** (e.g., the new entry is not created, discarded, or blocked through a blocking mechanism from being populated into the data warehouse **708**).

(68) In some embodiments, the data warehouse **708** is utilized to perform block sharing of data within the first volume **106**. In some embodiments, the data warehouse **708** is utilized to deduplicate data within the first volume **106**.

(69) In some embodiments of performing the overall resync, a send blocks operation is received such as by the first site **102** from the second site **104**. An entry within the data warehouse **708** is verified as whether the entry is a hole as part of writing to the data warehouse **708** via the send blocks operation. In some embodiments, the send blocks operation is for a virtual volume block number. A verification is performed to verify that the changed data of the virtual volume block number (e.g., changed data within the second volume **108** to resync to the first volume **106**) is not

being received in parallel with the send blocks operation.

(70) Hole punching is the process of removing spare blocks (holes) from copied blocks such as after a copy operation (e.g., a copy offload operation). An entry within the data warehouse **708** may be created for a virtual volume block number before the resync operation **702** transitions the virtual volume block number to a hole. Accordingly, a flag is set in an indirect (e.g., an indirect block/pointer) in the data warehouse **708** for the entry while setting the hole as part of the resync operation **702**. In some embodiments, a mapping may be applied to the data warehouse **078** for the virtual volume block number based upon the flag not being set. The mapping is not applied to the data warehouse **708** for the virtual block number based upon the flag being set (e.g., the data warehouse **708** is evaluated to determine whether the flag is set, and the mapping is not generated, is discarded, or is blocked from being populated within the data warehouse **708**). In some embodiments, a determination is made that the indirect of the data warehouse **708** is not allocated while removing a hole as part of the resync operation **702**. Accordingly, the indirect is allocated and the flag is set within the indirect.

(71) In some embodiments, a method is provided. The method includes replicating data from a first volume hosted at a first site as a primary volume to a second volume hosted at a second site as a secondary volume; in response to the first site experiencing a failure, promoting the second volume to be the primary volume and directing client I/O to the second volume for storing data modifications; in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and during the resync operation: logging database rebuild operations, for rebuilding a database file used to track blocks that are shared by files of at least one of the first volume or the second volume, into a log metafile; and implementing a background scanner to reconcile information within the log metafile for rebuilding the database file.

(72) In some embodiments, the method includes representing a data warehouse, used to track locations of blocks within at least one of the first volume or the second volume, utilizing both the database file and the log metafile; and utilizing both the database file and the log metafile to obtain a consistent state of at least one of the first volume or the second volume.

(73) In some embodiments, the method includes evaluating the log metafile to determine whether a block has been processed as part of the resync operation.

(74) In some embodiments, the method includes utilizing the log metafile and the database file to deduplicate blocks storing data of the first volume.

(75) In some embodiments, the method includes in response to the resync operation completing, demoting the second volume to become the secondary volume and promoting the first volume to become the primary volume.

(76) In some embodiments, the method includes maintaining a data warehouse for the secondary volume, wherein the data warehouse tracks blocks storing data of the secondary volume.

(77) In some embodiments, virtual volume block number mappings are logged into the log metafile.

(78) In some embodiments, the method includes populating a data warehouse maintained at the second site for the secondary volume with virtual volume block numbers used by the first volume.

(79) In some embodiments, the method includes populating a data warehouse maintained at the second site for the secondary volume with mappings between virtual volume block numbers, used by the first volume for storing the data, to locations used by the second volume for storing corresponding data.

(80) In some embodiments, a computing device is provided. The computing device includes a memory storing instructions and a processor coupled to the memory, the processor configured to execute the instructions to perform operations. The operations include replicating data from a first volume hosted at a first site as a primary volume to a second volume hosted at a second site as a secondary volume; in response to the first site experiencing a failure, promoting the second volume

to be the primary volume and directing client I/O to the second volume for storing data modifications; in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and during the resync operation: logging database rebuild operations, for rebuilding a database file used to track blocks that are shared by files of at least one of the first volume or the second volume, into a log metafile; and implementing a background scanner to reconcile information within the log metafile for rebuilding the database file.

(81) In some embodiments, the operations include logging virtual volume block numbers from the second site into the log metafile indexed by the virtual volume block numbers and mapped to virtual volume block numbers used by the first site.

(82) In some embodiments, the operations include logging virtual volume block numbers from the second site into the log metafile indexed by the virtual volume block numbers, wherein the indexes are at fixed offsets and a maximum number of indices is derived from a largest virtual volume block number from the second site.

(83) In some embodiments, the operations include specifying that the resync operation is complete in response to virtual volume block numbers from the second site being logged into the log metafile.

(84) In some embodiments, the operations include asynchronously processing map entries within the log metafile, wherein for each map entry from the second volume to the first volume, a block share is performed between a file block number within a data warehouse for the second volume to a container file block number associated with the first volume.

(85) In some embodiments, non-transitory machine readable medium is provided. The non-transitory machine readable medium comprises instructions for performing a method, which when executed by a machine, causes the machine to perform operations. The operations include replicating data from a first volume hosted at a first site as a primary volume to a second volume hosted at a second site as a secondary volume; in response to the first site experiencing a failure, promoting the second volume to be the primary volume and directing client I/O to the second volume for storing data modifications; in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and during the resync operation: logging database rebuild operations, for rebuilding a database file used to track blocks that are shared by files of at least one of the first volume or the second volume, into a log metafile; and implementing a background scanner to reconcile information within the log metafile for rebuilding the database file.

(86) In some embodiments, the operations include asynchronously processing map entries within the log metafile, wherein for each map entry from the second volume to the first volume, a block share is performed between a file block number within a data warehouse for the second volume to a container file block number associated with the first volume.

(87) In some embodiments, the operations include populating a data warehouse maintained at the second site for the secondary volume with virtual volume block numbers used by the first volume.

(88) In some embodiments, the operations include populating a data warehouse maintained at the second site for the secondary volume with mappings between virtual volume block numbers, used by the first volume for storing the data, to locations used by the second volume for storing corresponding data.

(89) In some embodiments, the operations include logging virtual volume block numbers from the second site into the log metafile indexed by the virtual volume block numbers and mapped to virtual volume block numbers used by the first site.

(90) In some embodiments, the operations include logging virtual volume block numbers from the second site into the log metafile indexed by the virtual volume block numbers, wherein the indexes are at fixed offsets and a maximum number of indices is derived from a largest virtual volume

block number from the second site.

(91) In some embodiments, a method is provided. The method includes in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and during the resync operation: evaluating a data warehouse maintained at the second site to identifying mappings between virtual volume block numbers used by the first volume to virtual volume block numbers used by the second volume; generating reverse mappings, mapping the virtual volume block numbers used by the second volume to the virtual volume block numbers used by the first volume, based upon the mappings from the data warehouse; and for each block being sent from the second site to the first site, generating and sending a command, including data of a block and a reverse mapping for the block, to the first site for updating the first volume.

(92) In some embodiments, the method includes in response to receiving the command at the first site, performing a block sharing operation to share the block based upon the reverse mapping embedded within the command.

(93) In some embodiments, the method includes constructing the command as a replication operation that describes an inode number, an offset, and a virtual volume block number of the second volume.

(94) In some embodiments, the method includes constructing, during the resync operation, a replication operation that describes an inode number, an offset, virtual volume block numbers of the second volume, and corresponding virtual volume block numbers of the first volume.

(95) In some embodiments, the method includes utilizing the reverse mappings within commands transmitted from the second site to the first site during the resync operation to perform deduplication for the first volume.

(96) In some embodiments, the method includes utilizing the reverse mappings within commands transmitted from the second site to the first site during the resync operation to preserve, at the first site, block sharing performed at the second site.

(97) In some embodiments, the method includes in response to determining that the data warehouse lacks a mapping for a virtual volume block number used by the second volume, transmitting a push name data command to the first site for performing a share operation.

(98) In some embodiments, the method includes in response to determining that the data warehouse lacks a mapping for a virtual volume block number used by the second volume based upon the virtual volume block number not being used at the first site, transmitting a push name data command to the first site for performing a share operation.

(99) In some embodiments, the method includes in response to determining that the data warehouse lacks a mapping for a virtual volume block number used by the second volume based upon the virtual volume block number being newly allocated as part of an incremental transfer from the second site to the first site, transmitting a push name data command to the first site for performing a share operation.

(100) In some embodiments, a computing device is provided. The computing device includes a memory storing instructions and a processor coupled to the memory, the processor configured to execute the instructions to perform operations. The operations include in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and during the resync operation: evaluating a data warehouse maintained at the second site to identifying mappings

between virtual volume block numbers used by the first volume to virtual volume block numbers used by the second volume; generating reverse mappings, mapping the virtual volume block numbers used by the second volume to the virtual volume block numbers used by the first volume, based upon the mappings from the data warehouse; and for each block being sent from the second site to the first site, generating and sending a command, including data of a block and a reverse mapping for the block, to the first site for updating the first volume.

(101) In some embodiments, the operations include in response to the resync operation completing, demoting the second volume to become the secondary volume and promoting the first volume to become the primary volume.

(102) In some embodiments, the operations include in response to receiving the command at the first site, performing a block sharing operation to share the block based upon the reverse mapping embedded within the command.

(103) In some embodiments, the operations include constructing the command as a replication operation that describes an inode number, an offset, and a virtual volume block number of the second volume.

(104) In some embodiments, the operations include constructing, during the resync operation, a replication operation that describes an inode number, an offset, virtual volume block numbers of the second volume, and corresponding virtual volume block numbers of the first volume.

(105) In some embodiments, the operations include utilizing the reverse mappings within commands transmitted from the second site to the first site during the resync operation to perform deduplication for the first volume.

(106) In some embodiments, non-transitory machine readable medium is provided. The non-transitory machine readable medium comprises instructions for performing a method, which when executed by a machine, causes the machine to perform operations. The operations include in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and during the resync operation: evaluating a data warehouse maintained at the second site to identifying mappings between virtual volume block numbers used by the first volume to virtual volume block numbers used by the second volume; generating reverse mappings, mapping the virtual volume block numbers used by the second volume to the virtual volume block numbers used by the first volume, based upon the mappings from the data warehouse; and for each block being sent from the second site to the first site, generating and sending a command, including data of a block and a reverse mapping for the block, to the first site for updating the first volume.

(107) In some embodiments, the operations include utilizing the reverse mappings within commands transmitted from the second site to the first site during the resync operation to preserve, at the first site, block sharing performed at the second site.

(108) In some embodiments, the operations include in response to determining that the data warehouse lacks a mapping for a virtual volume block number used by the second volume, transmitting a push name data command to the first site for performing a share operation.

(109) In some embodiments, the operations include in response to determining that the data warehouse lacks a mapping for a virtual volume block number used by the second volume based upon the virtual volume block number not being used at the first site, transmitting a push name data command to the first site for performing a share operation.

(110) In some embodiments, the operations include in response to determining that the data warehouse lacks a mapping for a virtual volume block number used by the second volume based upon the virtual volume block number being newly allocated as part of an incremental transfer from the second site to the first site, transmitting a push name data command to the first site for

performing a share operation.

(111) In some embodiments, a method is provided. The method includes in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; and in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number.

(112) In some embodiments, the method includes executing the background scanner to track a lifecycle of rebuilding the data warehouse by storing a changed vector of information with entries that include an inode, a file block number, a destination virtual volume block number, and size information; and retrieving source virtual volume block numbers for an entry within the changed vector of information to determine whether a new mapping is to be created within the data warehouse.

(113) In some embodiments, the method includes in response to determining that the data warehouse comprises the mapping between the changed virtual volume block number and the source virtual volume block number, retaining the mapping within the data warehouse.

(114) In some embodiments, the method includes in response to determining that there is a consistency group geometry difference between the first volume and the second volume resulting a lack of a one to one mapping between source and destination virtual volume block numbers, refraining from populating the entry within the data warehouse.

(115) In some embodiments, the method includes in response to determining that there is a compression difference between the first volume and the second volume resulting a lack of a one to one mapping between source and destination virtual volume block numbers, refraining from populating the entry within the data warehouse.

(116) In some embodiments, the method includes executing the snapshot difference scanner to generate a baseline difference by determining a difference between a null snapshot and a common snapshot between the first site and the second site.

(117) In some embodiments, the method includes executing the snapshot difference scanner to generate a baseline difference by determining a difference between a null snapshot and a public inofile of a common snapshot between the first site and the second site.

(118) In some embodiments, the method includes generating, by the snapshot difference scanner, a baseline difference by determining a difference between a null snapshot and a inofile of a common snapshot between the first site and the second site; issuing the baseline difference for each changed inode identified for the second volume to determine a changed vector of information; and retrieving source virtual volume block numbers for an entry within the changed vector of information to determine whether a new mapping is to be created within the data warehouse.

(119) In some embodiments, the method includes issuing a baseline difference for each changed inode identified for the second volume to determine a changed vector of information with entries that include an inode, a file block number, a destination virtual volume block number, and size information; and retrieving source virtual volume block numbers for an entry within the changed vector of information to determine whether a new mapping is to be created within the data warehouse.

(120) In some embodiments, the method includes utilizing the data warehouse to perform block

sharing of data within the first volume.

(121) In some embodiments, a computing device is provided. The computing device includes a memory storing instructions and a processor coupled to the memory, the processor configured to execute the instructions to perform operations. The operations include in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; and in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number.

(122) In some embodiments, the operations include utilizing the data warehouse to deduplicate data within the first volume.

(123) In some embodiments, the operations include receiving a send blocks operation; and verifying that an entry within the data warehouse is a hole as part of writing to the data warehouse via the send blocks operation.

(124) In some embodiments, the operations include receiving a send blocks operation for a virtual volume block number; and verifying that changed data of the virtual volume block number is not being received in parallel with the send blocks operation.

(125) In some embodiments, the operations include in response to creating an entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual volume block number to a hole, setting a flag within an indirect in the data warehouse for each entry while setting the hole as part of the resync operation.

(126) In some embodiments, non-transitory machine readable medium is provided. The non-transitory machine readable medium comprises instructions for performing a method, which when executed by a machine, causes the machine to perform operations. The operations include in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; and in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number.

(127) In some embodiments, the operations include in response to creating an entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual volume block number to a hole, setting a flag within an indirect in the data warehouse for one or more entries while setting the hole as part of the resync operation; and applying a mapping to the data warehouse for the virtual volume block number based upon the flag not being set.

(128) In some embodiments, the operations include in response to creating an entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual

volume block number to a hole, setting a flag within an indirect in the data warehouse for one or more entries while setting the hole as part of the resync operation; and refraining from applying a mapping to the data warehouse for the virtual volume block number based upon the flag being set. (129) In some embodiments, the operations include in response to determining that an indirect of the data warehouse is not allocated while removing a hole as part of the resync operation, allocating the indirect; and setting a flag within the indirect.

(130) In some embodiments, the operations include utilizing the data warehouse to perform block sharing of data within the first volume.

(131) Referring to FIG. 8, a node **800** (also referred to as a storage node) in this particular example includes processor(s) **801**, a memory **802**, a network adapter **804**, a cluster access adapter **806**, and a storage adapter **808** interconnected by a system bus **810**. In other examples, the node **800** comprises a virtual machine, such as a virtual storage machine.

(132) The node **800** also includes a storage operating system **812** installed in the memory **802** that can, for example, implement a RAID data loss protection and recovery scheme to optimize reconstruction of data of a failed disk or drive in an array, along with other functionality such as deduplication, compression, snapshot creation, data mirroring, synchronous replication, asynchronous replication, encryption, etc.

(133) The network adapter **804** in this example includes the mechanical, electrical and signaling circuitry needed to connect the node **800** to one or more of the client devices over network connections, which may comprise, among other things, a point-to-point connection or a shared medium, such as a local area network. In some examples, the network adapter **804** further communicates (e.g., using Transmission Control Protocol/Internet Protocol (TCP/IP)) via a cluster fabric and/or another network (e.g., a WAN (Wide Area Network)) (not shown) with storage devices of a distributed storage system to process storage operations associated with data stored thereon.

(134) The storage adapter **808** cooperates with the storage operating system **812** executing on the node **800** to access information requested by one of the client devices (e.g., to access data on a data storage device managed by a network storage controller). The information may be stored on any type of attached array of writeable media such as magnetic disk drives, flash memory, and/or any other similar media adapted to store information.

(135) In exemplary data storage devices, information can be stored in data blocks on disks. The storage adapter **808** can include I/O interface circuitry that couples to the disks over an I/O interconnect arrangement, such as a storage area network (SAN) protocol (e.g., Small Computer System Interface (SCSI), Internet SCSI (iSCSI), hyperSCSI, Fiber Channel Protocol (FCP)). The information is retrieved by the storage adapter **808** and, if necessary, processed by the processor(s) **801** (or the storage adapter **808** itself) prior to being forwarded over the system bus **810** to the network adapter **804** (and/or the cluster access adapter **806** if sending to another node computing device in the cluster) where the information is formatted into a data packet and returned to a requesting one of the client devices and/or sent to another node computing device attached via a cluster fabric. In some examples, a storage driver **814** in the memory **802** interfaces with the storage adapter to facilitate interactions with the data storage devices.

(136) The storage operating system **812** can also manage communications for the node **800** among other devices that may be in a clustered network, such as attached to the cluster fabric. Thus, the node **800** can respond to client device requests to manage data on one of the data storage devices or storage devices of the distributed storage system in accordance with the client device requests.

(137) The node **800** may implement a resync component **820** configured to perform the techniques described herein such as in relation to FIGS. 1A-1C and FIGS. 2-7. For example, the resync component **820** may perform the resync operation **130**, generate a data warehouse rebuild operation **132**, rebuild a data warehouse **134**, implement a background scanner **310**, generate commands **506**, implement a background scanner **706**, perform the method **200** of FIG. 2, perform the method **400**

of FIG. 4, and/or perform the method **600** of FIG. 6.

(138) In the example node **800**, memory **802** can include storage locations that are addressable by the processor(s) **801** and adapters **804**, **806**, and **808** for storing related software application code and data structures. The processor(s) **801** and adapters **804**, **806**, and **808** may, for example, include processing elements and/or logic circuitry configured to execute the software code and manipulate the data structures.

(139) The storage operating system **812**, portions of which are typically resident in the memory **802** and executed by the processor(s) **801**, invokes storage operations in support of a file service implemented by the node **800**. Other processing and memory mechanisms, including various computer readable media, may be used for storing and/or executing application instructions pertaining to the techniques described and illustrated herein.

(140) The examples of the technology described and illustrated herein may be embodied as one or more non-transitory computer or machine readable media, such as the memory **802**, having machine or processor-executable instructions stored thereon for one or more aspects of the present technology, which when executed by processor(s), such as processor(s) **801**, cause the processor(s) to carry out the steps necessary to implement the methods of this technology, as described and illustrated with the examples herein. In some examples, the executable instructions are configured to perform one or more steps of a method described and illustrated later.

(141) Still another embodiment involves a computer-readable medium **900** comprising processor-executable instructions configured to implement one or more of the techniques presented herein. An example embodiment of a computer-readable medium or a computer-readable device that is devised in these ways is illustrated in FIG. 9, wherein the implementation comprises a computer-readable medium **908**, such as a compact disc-recordable (CD-R), a digital versatile disc-recordable (DVD-R), flash drive, a platter of a hard disk drive, etc., on which is encoded computer-readable data **906**. This computer-readable data **906**, such as binary data comprising at least one of a zero or a one, in turn comprises processor-executable computer instructions **904** configured to operate according to one or more of the principles set forth herein. In some embodiments, the processor-executable computer instructions **904** are configured to perform a method **902** such as method **200** of FIG. 2, method **400** of FIG. 4, and/or method **600** of FIG. 6. In some embodiments, the processor-executable computer instructions **904** are configured to implement a system such as system **300** of FIG. 3, system **500** of FIG. 5, and/or system **700** of FIG. 7. Many such computer-readable media are contemplated to operate in accordance with the techniques presented herein.

(142) In an embodiment, the described methods and/or their equivalents may be implemented with computer executable instructions. Thus, in an embodiment, a non-transitory computer readable/storage medium is configured with stored computer executable instructions of an algorithm/executable application that when executed by a machine(s) cause the machine(s) (and/or associated components) to perform the method. Example machines include but are not limited to a processor, a computer, a server operating in a cloud computing system, a server configured in a Software as a Service (SaaS) architecture, a smart phone, and so on. In an embodiment, a computing device is implemented with one or more executable algorithms that are configured to perform any of the disclosed methods.

(143) It will be appreciated that processes, architectures and/or procedures described herein can be implemented in hardware, firmware and/or software. It will also be appreciated that the provisions set forth herein may apply to any type of special-purpose computer (e.g., file host, storage server and/or storage serving appliance) and/or general-purpose computer, including a standalone computer or portion thereof, embodied as or including a storage system. Moreover, the teachings herein can be configured to a variety of storage system architectures including, but not limited to, a network-attached storage environment and/or a storage area network and disk assembly directly attached to a client or host computer. Storage system should therefore be taken broadly to include such arrangements in addition to any subsystems configured to perform a storage function and

associated with other equipment or systems.

(144) In some embodiments, methods described and/or illustrated in this disclosure may be realized in whole or in part on computer-readable media. Computer readable media can include processor-executable instructions configured to implement one or more of the methods presented herein, and may include any mechanism for storing this data that can be thereafter read by a computer system. Examples of computer readable media include (hard) drives (e.g., accessible via network attached storage (NAS)), Storage Area Networks (SAN), volatile and non-volatile memory, such as read-only memory (ROM), random-access memory (RAM), electrically erasable programmable read-only memory (EEPROM) and/or flash memory, compact disk read only memory (CD-ROM) s, CD-Rs, compact disk re-writeable (CD-RW) s, DVDs, cassettes, magnetic tape, magnetic disk storage, optical or non-optical data storage devices and/or any other medium which can be used to store data.

(145) Some examples of the claimed subject matter have been described with reference to the drawings, where like reference numerals are generally used to refer to like elements throughout. In the description, for purposes of explanation, numerous specific details are set forth in order to provide an understanding of the claimed subject matter. It may be evident, however, that the claimed subject matter may be practiced without these specific details. Nothing in this detailed description is admitted as prior art.

(146) Although the subject matter has been described in language specific to structural features or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as example forms of implementing at least some of the claims.

(147) Various operations of embodiments are provided herein. The order in which some or all of the operations are described should not be construed to imply that these operations are necessarily order dependent. Alternative ordering will be appreciated given the benefit of this description. Further, it will be understood that not all operations are necessarily present in each embodiment provided herein. Also, it will be understood that not all operations are necessary in some embodiments.

(148) Furthermore, the claimed subject matter is implemented as a method, apparatus, or article of manufacture using standard application or engineering techniques to produce software, firmware, hardware, or any combination thereof to control a computer to implement the disclosed subject matter. The term “article of manufacture” as used herein is intended to encompass a computer application accessible from any computer-readable device, carrier, or media. Of course, many modifications may be made to this configuration without departing from the scope or spirit of the claimed subject matter.

(149) As used in this application, the terms “component”, “module,” “system”, “interface”, and the like are generally intended to refer to a computer-related entity, either hardware, a combination of hardware and software, software, or software in execution. For example, a component includes a process running on a processor, a processor, an object, an executable, a thread of execution, an application, or a computer. By way of illustration, both an application running on a controller and the controller can be a component. One or more components residing within a process or thread of execution and a component may be localized on one computer or distributed between two or more computers.

(150) Moreover, “exemplary” is used herein to mean serving as an example, instance, illustration, etc., and not necessarily as advantageous. As used in this application, “or” is intended to mean an inclusive “or” rather than an exclusive “or”. In addition, “a” and “an” as used in this application are generally be construed to mean “one or more” unless specified otherwise or clear from context to be directed to a singular form. Also, at least one of A and B and/or the like generally means A or B and/or both A and B. Furthermore, to the extent that “includes”, “having”, “has”, “with”, or

variants thereof are used, such terms are intended to be inclusive in a manner similar to the term “comprising”.

(151) Many modifications may be made to the instant disclosure without departing from the scope or spirit of the claimed subject matter. Unless specified otherwise, “first,” “second,” or the like are not intended to imply a temporal aspect, a spatial aspect, an ordering, etc. Rather, such terms are merely used as identifiers, names, etc. for features, elements, items, etc. For example, a first set of information and a second set of information generally correspond to set of information A and set of information B or two different or two identical sets of information or the same set of information.

(152) Also, although the disclosure has been shown and described with respect to one or more implementations, equivalent alterations and modifications will occur to others skilled in the art based upon a reading and understanding of this specification and the annexed drawings. The disclosure includes all such modifications and alterations and is limited only by the scope of the following claims. In particular regard to the various functions performed by the above described components (e.g., elements, resources, etc.), the terms used to describe such components are intended to correspond, unless otherwise indicated, to any component which performs the specified function of the described component (e.g., that is functionally equivalent), even though not structurally equivalent to the disclosed structure. In addition, while a particular feature of the disclosure may have been disclosed with respect to only one of several implementations, such feature may be combined with one or more other features of the other implementations as may be desired and advantageous for any given or particular application.

Claims

1. A method, comprising: in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: tracking the rebuild of the data warehouse using a changed vector of changes, to the second volume since the failure, identified based upon changed inode differences; for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the changed vector to determine whether a new mapping is to be created within the data warehouse; and in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the new mapping between the changed virtual volume block number and the source virtual volume block number.

2. The method of claim 1, comprising: generate a baseline difference by determining a difference between a snapshot and a common snapshot between the first site and the second site.

3. The method of claim 1, comprising: in response to determining that the data warehouse comprises the mapping between the changed virtual volume block number and the source virtual volume block number, retaining the mapping within the data warehouse.

4. The method of claim 1, comprising: in response to determining that there is a geometry difference between the first volume and the second volume resulting in a lack of a one to one mapping between source and destination virtual volume block numbers, refraining from populating the entry within the data warehouse.

5. The method of claim 1, comprising: in response to determining that there is a compression difference between the first volume and the second volume resulting in a lack of a one to one mapping between source and destination virtual volume block numbers, refraining from populating

the entry within the data warehouse.

6. The method of claim 1, comprising: executing the snapshot difference scanner to generate a baseline difference by determining a difference between a null snapshot and a common snapshot between the first site and the second site.

7. The method of claim 1, comprising: executing the snapshot difference scanner to generate a baseline difference by determining a difference between a null snapshot and a public inofile of a common snapshot between the first site and the second site.

8. The method of claim 1, comprising: generating, by the snapshot difference scanner, a baseline difference by determining a difference between a null snapshot and a inofile of a common snapshot between the first site and the second site; issuing the baseline difference for each changed inode identified for the second volume to determine the changed vector of information.

9. The method of claim 1, comprising: issuing a baseline difference for each changed inode identified for the second volume to determine the changed vector of information.

10. The method of claim 1, comprising: utilizing the data warehouse to perform block sharing of data within the first volume.

11. A computing device comprising: a memory storing instructions; and a processor coupled to the memory, the processor configured to execute the instructions to perform operations comprising: in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; and in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number; and in response to determining that there is a geometry difference between the first volume and the second volume resulting in a lack of a one to one mapping between source and destination virtual volume block numbers, refraining from populating the entry within the data warehouse.

12. The computing device of claim 11, wherein the operations comprise: utilizing the data warehouse to deduplicate data within the first volume.

13. The computing device of claim 11, wherein the operations comprise: receiving a send blocks operation; and verifying that an entry within the data warehouse is a hole as part of writing to the data warehouse via the send blocks operation.

14. The computing device of claim 11, wherein the operations comprise: receiving a send blocks operation for a virtual volume block number; and verifying that changed data of the virtual volume block number is not being received in parallel with the send blocks operation.

15. The computing device of claim 11, wherein the operations comprise: in response to creating a new entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual volume block number to a hole, setting a flag within an indirect in the data warehouse for each entry while setting the hole as part of the resync operation.

16. A non-transitory machine readable medium comprising instructions for performing a method, which when executed by a machine, causes the machine to perform operations comprising: in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications

from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number; and in response to determining that there is a compression difference between the first volume and the second volume resulting in a lack of a one to one mapping between source and destination virtual volume block numbers, retraining from populating the entry within the data warehouse.

17. The non-transitory machine readable medium of claim 16, wherein the operations comprise: in response to creating a new entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual volume block number to a hole, setting a flag within an indirect in the data warehouse for one or more entries while setting the hole as part of the resync operation; and applying a mapping to the data warehouse for the virtual volume block number based upon the flag not being set.

18. The non-transitory machine readable medium of claim 16, wherein the operations comprise: in response to creating a new entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual volume block number to a hole, setting a flag within an indirect in the data warehouse for one or more entries while setting the hole as part of the resync operation; and refraining from applying a mapping to the data warehouse for the virtual volume block number based upon the flag being set.

19. The non-transitory machine readable medium of claim 16, wherein the operations comprise: in response to determining that an indirect of the data warehouse is not allocated while removing a hole as part of the resync operation, allocating the indirect; and setting a flag within the indirect.

20. The non-transitory machine readable medium of claim 16, wherein the operations comprise: utilizing the data warehouse to perform block sharing of data within the first volume.

21. A method, comprising: in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number; generating, by the snapshot difference scanner, a baseline difference by determining a difference between a null snapshot and a inofile of a common snapshot between the first site and the second site; issuing the baseline difference for each changed inode identified for the second volume to determine a changed vector of information; and retrieving source virtual volume block numbers for a changed vector entry within the changed vector of information to determine whether a new mapping is to be created within the data warehouse.

22. The method of claim 21, comprising: utilizing the data warehouse to perform block sharing of data within the first volume.

23. The method of claim 21, comprising: utilizing the data warehouse to deduplicate data within the first volume.

24. A method, comprising: in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number; issuing a baseline difference for each changed inode identified for the second volume to determine a changed vector of information with entries that include an inode, a file block number, a destination virtual volume block number, and size information; and retrieving source virtual volume block numbers for a changed vector entry within the changed vector of information to determine whether a new mapping is to be created within the data warehouse.

25. The method of claim 24, comprising: in response to determining that an indirect of the data warehouse is not allocated while removing a hole as part of the resync operation, allocating the indirect; and setting a flag within the indirect.

26. The method of claim 24, comprising: utilizing the data warehouse to deduplicate data within the first volume.

27. A non-transitory machine readable medium comprising instructions for performing a method, which when executed by a machine, causes the machine to perform operations comprising: in response to a first site, hosting a first volume as a primary volume, experiencing a failure, promoting a second volume, hosted by a second site as a secondary volume, to be the primary volume and directing client I/O to the second volume for storing data modifications; and in response to the first site recovering, initiating a resync operation to transfer the data modifications from the second volume to the first volume for re-promoting the first volume to be the primary volume; and executing a background scanner to rebuild a data warehouse at the first site by: for a changed virtual volume block number of the second volume identified by a snapshot difference scanner using one or more snapshots, obtaining a source virtual volume block number from the first site; in response to determining that the data warehouse does not comprise a mapping between the changed virtual volume block number and the source virtual volume block number, creating an entry within the data warehouse to create the mapping between the changed virtual volume block number and the source virtual volume block number; and in response to creating an entry within the data warehouse for a virtual volume block number before the resync operation transitions the virtual volume block number to a hole, setting a flag within an indirect in the data warehouse for each entry while setting the hole as part of the resync operation.

28. The non-transitory machine readable medium of claim 17, wherein the operations comprise: receiving a send blocks operation; and verifying that an entry within the data warehouse is the hole as part of writing to the data warehouse via the send blocks operation.

29. The non-transitory machine readable medium of claim 17, wherein the operations comprise: receiving a send blocks operation for a virtual volume block number; and verifying that changed data of the virtual volume block number is not being received in parallel with the send blocks operation.
