

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12386722
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Moser; Herwig et al.

---

### Method and system for real-time correlation of disparate sources of topological information to create a unified topological model of a distributed information system

---

#### Abstract

A system and method is disclosed for the combined analysis of transaction execution monitoring data and a topology model created from infrastructure monitoring data of computing systems involved in the execution of the monitored transactions. Monitored communication activities of transactions are analyzed to identify intermediate processing nodes between sender and receiver side and to enrich transaction monitoring data with data describing those intermediate processing nodes. The topology model may also be improved by the combined analysis, as functionality and services provided by elements of the topology model may be derived by the involvement of those elements in the execution of monitored transactions. The result of the combined analysis is used by an automated anomaly detection and causality estimation system. The combined analysis may also reveal entities of a monitored environment that are used by transaction executions but which are not monitored.

---

**Inventors:** Moser; Herwig (Freistadt, AT), Kopp; Michael (Piberbach, AT), Ambichl; Ernst (Altenberg, AT)

**Applicant:** Dynatrace LLC (Waltham, MA)

**Family ID:** 65724153

**Assignee:** Dynatrace LLC (Waltham, MA)

**Appl. No.:** 17/871282

**Filed:** July 22, 2022

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20230008791 A1	Jan. 12, 2023

## Related U.S. Application Data

continuation parent-doc US 16276710 20190215 US 11442836 child-doc US 17871282  
us-provisional-application US 62634393 20180223

---

## Publication Classification

**Int. Cl.:** **G06F11/34** (20060101); **G06F11/07** (20060101); **G06F11/30** (20060101); **G06N5/022** (20230101); **G06N5/04** (20230101)

**U.S. Cl.:**

**CPC** **G06F11/3447** (20130101); **G06F11/079** (20130101); **G06F11/3048** (20130101); **G06F11/3089** (20130101); **G06F11/3461** (20130101); **G06F11/3495** (20130101); **G06N5/022** (20130101); **G06N5/04** (20130101);

## Field of Classification Search

**USPC:** None

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9557879	12/2016	Wang et al.	N/A	N/A
2005/0251371	12/2004	Chagoly et al.	N/A	N/A
2009/0144409	12/2008	Dickerson et al.	N/A	N/A
2010/0296411	12/2009	Chambers	N/A	N/A
2012/0304172	12/2011	Greifeneder et al.	N/A	N/A
2013/0185643	12/2012	Greifeneder et al.	N/A	N/A
2014/0136693	12/2013	Greifeneder et al.	N/A	N/A
2014/0297846	12/2013	Hoja et al.	N/A	N/A
2015/0032884	12/2014	Greifeneder et al.	N/A	N/A
2015/0052441	12/2014	Degioanni	N/A	N/A
2016/0105350	12/2015	Greifeneder et al.	N/A	N/A
2016/0283345	12/2015	Gounares et al.	N/A	N/A
2017/0075749	12/2016	Ambichl et al.	N/A	N/A
2018/0075152	12/2017	Zhang	N/A	G06F 9/455
2018/0217879	12/2017	Garcia et al.	N/A	N/A

---

*Primary Examiner:* Chen; Qing

*Attorney, Agent or Firm:* Harness, Dickey & Pierce, P.L.C.

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 16/276,710 (issued as U.S. Pat. No. 11,442,836), filed on Feb. 15, 2019, which claims the benefit of U.S. Provisional Application No. 62/634,393, filed on Feb. 23, 2018. The entire disclosures of the above applications are incorporated herein by reference.

## FIELD

(1) The invention generally relates to the monitoring of transaction execution environments in combination with the monitoring of the executions of transactions by these environments and specifically to the combined analysis of environment and transaction monitoring data to improve the knowledge about the transaction execution environment and to enrich the transaction monitoring data with context data extracted from execution environment data.

## BACKGROUND

(2) Modern application monitoring systems employ a holistic monitoring approach, which combines monitoring the structure of application infrastructure, monitoring the resource usage of different elements of the application infrastructure and monitoring of transaction executions on the application infrastructure. The combined monitoring data is used to create a topological model of the application infrastructure and to map resource utilization data and transaction execution monitoring data to the entities of this topological model.

(3) Those modern monitoring systems typically use agents that are deployed to components of monitored environments like host computers/operating systems, processes and virtualization/cloud management entities. Those agents provide the topological, infrastructure monitoring and transaction monitoring data which is used by the monitoring system to create a holistic model of the monitored environment which is the basis for higher-level monitoring tasks like anomaly detection, estimation of causal relationships between anomalies and the identification of root cause anomalies.

(4) To get a most realistic model of the monitored environment, it is desired to deploy an agent to all its components, but technological or organizational impediments may prevent this. As an example, some components of the monitored environment may consist of dedicated hardware appliances that do not provide interfaces for internal, agent-based or external monitoring. Some components of the environment may not be controlled by the owner/operator of the monitoring system, therefore installing agents to those components may be possible from a technological point of view but impossible due to organizational obstacles. Finally, the operator of the monitoring system may simply have forgotten to deploy agents to some components.

(5) Not monitored components of the monitored environment may lead to an incomplete model of the monitored environment which may, in turn, lead to incorrect or misleading results of the higher-level monitoring tasks like causality and root cause estimation.

(6) Transaction monitoring data may contain data describing communication activities between two endpoints like e.g. a client and a server process. It may, in addition to data describing the topological location of the endpoints and measurement data describing the performance of the monitored communication, also contain data describing the path of the communication, e.g. in form of data identifying intermediate components between the endpoints. In addition, monitoring data provided by virtualization monitoring agents may contain data describing running virtualized computer systems, regardless of an operating system agent running on the virtualized computer system.

(7) This additional data provided by transaction and virtualization monitoring agents represents additional data describing the topology of monitored environments that may reveal the existence of not yet monitored or unmonitable (e.g. due to a lack of agent interfaces) components of the monitored environment. This additional data could be used to enrich both the topological model of the monitored environment and the recorded transaction trace data. The benefits of such enriched topology and transaction data include more accurate and precise results of higher-level monitoring

tasks like causality and root cause estimation.

(8) Consequently, a method and system are required in the art for a combined analysis of transaction monitoring and topology model data to infer additional topology model and transaction execution data to improve both the topology model and the transaction monitoring data.

(9) This section provides background information related to the present disclosure which is not necessarily prior art.

## SUMMARY

(10) This section provides a general summary of the disclosure, and is not a comprehensive disclosure of its full scope or all of its features.

(11) The disclosed technology is directed to the combined and interlaced analysis of topology and transaction monitoring data to infer the existence of not monitored components that are used by monitored transactions and to enrich both the topological model of a monitored environment and transaction trace data describing transaction executions performed by the monitored environment with data of the inferred, not monitored components. The improved topology model and transaction data may be used as the basis for an enhanced detection of abnormal performance or error behavior and for an improved localization of those detected anomalies in the topology model, which in turn leads to a more accurate and precise causality and root cause estimation for multiple detected and localized anomalies.

(12) Various agents are deployed to the components of a monitored environment, those agents may be of different types according to the type of component they are deployed to and according to the type of monitoring data they provide. Example agent types are cloud and virtualization agents for the monitoring of virtualization and cloud based environments that e.g. provide topology and resource usage data of virtualization and cloud related components, like hypervisor components hosting multiple virtualized computer systems, individual virtualized computer systems or cloud infrastructure components like load balancers or cloud endpoint services like relational database or file storage services, operating system agents for the monitoring of operating systems running on virtualized or real computer systems that e.g. provide topology and resource usage data of monitored operating systems, of processes running on those operating systems, data describing communication activities of monitored processes and of container based virtualization and process isolation components executed on the monitored operating systems, and process and transaction monitoring agents deployed to individual processes running in the monitored environment, that provide topology, resource usage and transaction tracing data from individual processes.

(13) The agents send gathered topology, resource usage, and transaction trace data to a monitoring server or an interconnected cluster of monitoring servers.

(14) The monitoring server uses the received topology data to incrementally build a multidimensional topology model of the monitored environment that integrates infrastructure, virtualization and transaction execution and service-related aspects of the monitored environment.

(15) Received transaction trace data is used to incrementally create end-to-end transaction monitoring data describing individual transaction executions. Portions of the created end-to-end transaction monitoring data corresponding to services provided by individual processes are correlated with topology entities representing those transactions to also integrate transaction monitoring data with the topology model of the monitored environment.

(16) The monitoring server may, in addition, perform a combined analysis of received transaction monitoring and topology monitoring data to infer the existence of components of the monitored environment that are not directly reported by deployed agents, e.g. because the nature of those components does not allow to deploy an agent to them.

(17) Both transaction monitoring data and topology model may be enriched with the inferred components.

(18) Variant embodiments may maintain a topology model based only on topology monitoring data reported by agents and a separate topology model that in addition contains inferred components to

avoid that already inferred components may be used as the basis for the inference of the existence of further not reported components. Further areas of applicability will become apparent from the description provided herein. The description and specific examples in this summary are intended for purposes of illustration only and are not intended to limit the scope of the present disclosure.

---

## Description

### DRAWINGS

- (1) The drawings described herein are for illustrative purposes only of selected embodiments and not all possible implementations, and are not intended to limit the scope of the present disclosure.
- (2) FIG. 1a provides a block diagram of a monitoring system monitoring an application execution environment that is partly implemented using cloud computing components.
- (3) FIG. 1b shows a monitored application execution environment containing multiple proxy entities that transfer transaction communication between different application servers.
- (4) FIG. 1c shows a monitored application execution environment with applications involved in the execution of monitored transactions that run in container environments.
- (5) FIG. 2a-c provide flow charts that conceptually describe the agent side capturing and sending of transaction execution environment or topology monitoring data.
- (6) FIG. 3a-c show flow charts that describe the agent side capturing and sending of transaction execution monitoring data.
- (7) FIG. 4a-i show data records that may be used to transfer transaction and topology monitoring data.
- (8) FIG. 5a-f depict data records that may be used for the representation of transition execution and topology data on the monitoring server.
- (9) FIG. 6a-b provide flow charts that conceptually describe the processing of received topology monitoring data by the monitoring server.
- (10) FIG. 7a-b show flow charts describing the processing of received transaction monitoring data by the monitoring server.
- (11) FIG. 8 provides a flowchart showing an overview of the transaction/topology data inference process to improve both transaction and topology data.
- (12) FIG. 9 shows a flowchart describing the enrichment of transaction monitoring data using topology monitoring data of the topology entities that executed the transaction.
- (13) FIG. 10 shows a flowchart describing the processing of a transaction processing entity like a computer host system reported by transaction monitoring data that is not available in the topology model.
- (14) FIG. 11 provides a flowchart showing the processing of sender side endpoint data of a reported transaction communication to enrich both the topology model and the transaction monitoring data.
- (15) FIG. 12 shows a flowchart describing the processing of receiver side endpoint data of a transaction communication to identify intermediate traffic management nodes like proxies between sender and receiver and to enrich both transaction monitoring and topology data accordingly.
- (16) FIG. 13 shows a flowchart that describes the identification of process level data of previously identified intermediate traffic management nodes.
- (17) FIG. 14 shows a flowchart describing the process of analyzing server and receivers side endpoint data of a monitored transaction communication to identify further intermediate traffic management nodes.
- (18) FIG. 15 shows a flowchart that illustrates the usage of enriched transaction monitoring and topology data to identify anomalies transaction execution behavior and to further identify causal relationships and the root cause for such identified anomalies.
- (19) Corresponding reference numerals indicate corresponding parts throughout the several views

of the drawings.

## DETAILED DESCRIPTION

(20) Example embodiments will now be described more fully with reference to the accompanying drawings.

(21) Embodiments of the present technology are directed to the combined and interleaved analysis of topology and transaction monitoring data to infer the existence of not monitored components of a monitored environment and to either enrich topology and transaction monitoring data with data describing the inferred components or to notify the user of a monitoring system that the existence of unmonitored components of the monitored environment was detected.

(22) Topology and transaction monitoring data are typically provided by agents deployed to the components of a monitored environment. Different types of agents may be deployed to different types of monitored component. Those different types of agents may provide monitoring data describing different aspects of the component they are deployed to.

(23) As an example, cloud and virtualization environment agents may be configured to monitor virtualization components that provide virtualized computer systems, like hypervisor components, or they may monitor cloud utilization data like number and type of cloud computing systems currently used by a monitored environment, or they may monitor cloud-inherent infrastructure components, like load balancing or proxy services provided by the cloud environment to control the flow of requests and data. Further, cloud and virtualization environment agents may provide topological monitoring data e.g. describing the relationships between virtualization provider components and virtualized computer systems (e.g. which virtualized computer system is provided by which hypervisor instance), or they may provide monitoring data describing the data flow created by data flow components of a monitored cloud environment (e.g. data paths established by load balancing components from the load balancing component to various receivers of the balanced load). Cloud and virtualization agents may further monitor and report network topology related aspects of the monitored environment, like virtual private cloud configuration data which defines network accessibility and visibility rules of cloud components.

(24) Operating system agents (or OS agents) may be deployed to operating systems running on concrete or (cloud) virtualized computer systems and may provide resource availability and usage data of the monitored operating system. OS agents may further provide correlation data identifying the operating system instance they are deployed to which may also be used to identify the corresponding virtualized computer system running the operating system that is monitored and reported by a virtualization and cloud agent. The OS agent may further monitor resource utilization and communication activities of processes running on the monitored operating system and it may also provide topology data describing the monitored operating systems and the processes running on it. OS agents may in addition monitor container-based virtualization and process isolation performed by container management processes executed on monitored operating systems. Those container management processes may create isolated execution environments for processes. Isolated execution environments may e.g. simulate a separate operating system environment for a process. This simulated operating system environments may, in addition to a separated process execution environment, provide a configured view on the network environment of the hosting operating system, including a mapping of component internal to component external network identifiers, a restricted view on network components from inside the container or even separated multi-container networks, like overlay-networks provided by docker container virtualization environments.

(25) Process or transaction agents may be deployed to processes running on operating systems of the monitored environment and may provide fine-grained resource utilization, communication and transaction execution monitoring data. The process or transaction agents may create and report topology data describing services (e.g. HTTP services or remote method invocation services) provided by the monitored processes and they may provide transaction trace and monitoring data

that describe transaction executions performed by the monitored processes. Process or transaction agents run inside the process context of the monitored process and therefore have the same level of access to process internal data as other (application) code that is executed by the process.

Therefore, process or transaction agents may also be referred to as in-process agents.

(26) The provided transaction trace and monitoring data may, in addition, contain data identifying the services used by the monitored transactions.

(27) The topology data provided by the different agents is received by an individual monitoring server or a monitoring server cluster and may be used to incrementally create a topological model of the monitored environment. Resource usage and transaction trace data are also received by the monitoring server or monitoring server cluster. The received transaction trace data is processed to create end-to-end transaction trace data. Resource utilization and transaction execution performance measurement data extracted from received transaction trace data are analyzed to identify abnormal resource usage or transaction execution performance. The identified abnormal conditions are analyzed in view of the known topology model to identify causal dependencies between different abnormal conditions and to identify for multiple, causally related abnormal conditions, those abnormal conditions out of the causally related abnormal conditions with the highest probability of being the root cause of all identified causally related abnormal conditions.

(28) The individual or clustered monitoring servers may, before performing anomaly detection and root cause calculation, perform a combined analysis of received topology data and transaction trace data to identify evidence of components existing in the monitored environment for which no direct topology data is available, e.g. because no agent is deployed to monitor those components for various reasons. The evidence for those existing but not monitored components may be used to create a version of the topology model that also describes those not directly monitored components. Further, this evidence may be used to enrich end-to-end transaction trace data containing evidence that those not directly monitored components are used with additional data describing the usage of those not directly monitored components by the transaction. In addition a user or administrator of the monitoring system may be notified about the existence of those not directly monitored components together with a suggestion to add agents to those not directly monitored components.

(29) Referring now to FIG. 1a, which provides a block diagram of a monitoring system consisting of a monitoring server **101** that receives topology data **121** and transaction trace and monitoring data **120** from various types of agents deployed to components of a monitored environment.

(30) The exemplary monitored environment depicted in FIG. 1a consists of a not virtualized host computer system **130**, running a process that executes application 1 **131**. A transaction agent **150a** is deployed to application 1 and monitors transactions (e.g. T1 **140** and T2 **141**) executed by application 1. An OS agent **151a** is deployed to the operating system executed on host 1 **130** and monitors overall resource usage of the operating system executed on host 1. The OS agent **151a** also monitors resource usage and communication activity of processes running on the operating system, like the process running application 1 **131**. OS agent **151a** provides topology monitoring data describing and identifying the computing system host 1, processes executing on host 1, like the process **131** running application 1 **131**, and topology data describing communication activities of the processes running on host 1.

(31) Further, the monitored environment contains components running in a cloud/virtualization environment **132**, like a load balancing service **133** and virtualized, cloud computing systems host 2 **136** and host 3 **134**, running processes application 2 **137** and application 3 **135**. A cloud and virtualization monitoring agent **152** is installed in the cloud environment **132** which uses monitoring interfaces of the cloud management system. Those monitoring interfaces of the cloud management system are used to provide and control the cloud environment **132** to request and report data describing cloud components and their relationships, including communication and usage relationships. The cloud and virtualization monitoring agent **152** may be implemented as process running on a host computing system (virtualized or not virtualized), the host computing

system may run inside or outside the monitored cloud/virtualization environment, as long as the cloud and virtualization agent is able to access the monitoring interfaces of the cloud management systems. In some variant embodiments, the cloud and virtualization monitoring agent may be implemented as functionality of a part of the monitoring system components, like the monitoring server **101**.

(32) Further, the cloud and virtualization monitoring agent **152** may fetch and report network configuration data, like configuration data defining virtual private cloud environments that define accessibility and visibility of individual components of the cloud environment to form networks of isolated subsets of cloud components with a defined interface, e.g. in form of a load balancer configured to receive requests from outside the virtual private cloud and routing them to components of the virtual private cloud.

(33) Example interfaces that may be used by the cloud and virtualization agent are the Amazon CloudWatch® interface for AWS cloud environments as e.g. provided by Amazon.com, Inc. or Azure Monitor for Azure cloud environments as e.g. provided by the Microsoft Corporation.

(34) An OS agent **151b** may be deployed to host 3 **134**, which monitors **153** the process running application 3 **135**, and a transaction agent **150b** may be deployed to the process of application 3 **135** which monitors transactions executed by application 3.

(35) Host 4 **138**, which runs application 4 **139**, also belongs to the monitored environment. An OS agent **151c** is deployed to host 4, which monitors the process that runs application 4. In this case, no transaction agent **150** is deployed to the application process for one of various reasons, including the lack of availability of an agent for the application architecture used for application 4, or a monitoring configuration that excludes the deployment of a transaction agent to application 4. OS agent **151c** provides topology monitoring data describing and identifying the computing system host 4, processes executing on host 4, like the process **139** running application 4 **139**, and topology data describing communication activities of the processes running on host 4.

(36) The cloud and virtualization agent **152** monitors **154** the cloud and virtualization environment **132** and provides topology data describing the existence of the virtualized cloud hosts host 2 **136** and host 3 **134** and the load balancing service **133**. Further, the cloud and virtualization agent **152** may query configuration data of host 2 **136** and host 3 **134** describing the resources of the virtual computing system in terms of CPU, main memory, and persistent memory, and it may query configuration and capacity data of the load balancing service **133** in terms of maximal rate of processed requests and input/output mapping data describing one or more network endpoints (i.e. network address and port for TCP/IP based networks) at which the load balancing service expects incoming requests and corresponding mapped network endpoints to which the load balancing service distributes received requests. The topology data created by the cloud and virtualization data may further describe that the load balancing service forwards received requests to host 2 **136** and host 3 **134**.

(37) An OS agent **151b** is deployed to the virtualized cloud computing system host **134**, which monitors the operating system executed on host **134**. The OS agent **151b** also monitors **153** processes executed on the operating system, like the process running application 3 **135**. The topology data provided by the OS agent **151b** includes data describing and identifying the monitored operating system running on host 3 **134** and data identifying the host computing system host 3 **134**, where the identifying data for the host computing system provided by the OS agent contains data that may be used to identify the corresponding virtual computing system reported by the cloud and virtualization agent **152**. The OS agent further provides topology data describing the processes executed by the operating system of host 3, like the process running application 3 **135**, and data describing the communication activities of those processes.

(38) Transactions T1 **140** and T2 **142** are currently executed by the monitored environment, both transactions enter the system via application 1 running on host 1. Both transactions are initially processed by application 1 and then forwarded to the load balancing service **133** of the cloud



portion of the monitored environment. The load balancing service **133** decides, based on its configuration data, that transaction T1 **140** is forwarded to application 2 **137** running on host 2 **136** and that transaction T2 **142** is forwarded to application 3 **135** running on host 3 **134**. Processing of T1 **140** is finished on application 2. Processing of transaction T2 continues on application 3 **135**, which is executed on host 3 **134**, and is then forwarded to application 4 on host **138**, where processing of transaction T2 is finished.

(39) The transaction agents **150a** and **150b** deployed to processes **131** running application 1 **131** and application 3 **135** monitor the execution of transactions, like transactions T1 **140** and T2 **141** by the processes they are deployed to and provide transaction trace data fragments describing portions of the execution of transactions together with correlation data that may be used to identify and correlate transaction trace data fragments from different agents that describe different portions of the same transaction. Portions of transaction data fragments provided by the transaction agents that describe interprocess communication (e.g. for transaction T1 **140** the portion of T1 executed on application 1 that caused the execution of a portion of T1 on application 2 on host 2), contain data describing both communication endpoints as seen from the sender or receiver process.

(40) As an example, the portion of the transaction trace fragment describing communication between application 1 and application 2 as provided by transaction agent **150a** contains data describing the communication endpoint identifying application) **131** on host 1 **130** as sender endpoint, and the communication endpoint of the load balancing service **133** to which the communication message was sent as receiver endpoint, as the load balancing service acts as receiver of communication messages or requests for sender applications, even if it only forwards the communication message or request to a real receiver, like application 2 **137** on host 2 **136**. As another example, application 3 receives a request to continue the execution of transaction T2 sent from application 1 via the load balancing service **133**. Even if the request was originally sent by application 1, for the receiver application 3 the visible network endpoint from which the request was received is the load balancing service.

(41) Transaction trace data created by transaction agents **150a**, **150b** and **150c** contains, next to data describing the execution of transactions, topology localization data that identifies the component that executed a portion of a transaction (e.g. a process) with the corresponding topology entity describing this component. The topology localization data allows localizing the position of transaction executions within the topology model.

(42) The transaction agents send the topology localized transaction trace data **120** to a monitoring server **101**, which forwards it to a transaction trace data processor **102**. The transaction trace data processor correlates transaction trace data corresponding to the same monitored transaction to incrementally create end-to-end trace data **123** describing the whole monitored transaction.

(43) Topology data **121** as sent by OS agents and cloud and virtualization agents, is processed by a topology data processor component **104** of the monitoring server to create a topology model **105** of the monitored environment as reported by the various agents. The topology data processor may in addition process transaction trace data **120** and end-to-end transaction data to extract data describing services provided by processes involved in the execution of those transactions, to extract service call relationships out of the processed transaction data and to enrich the topology model **105** with this service data (not shown).

(44) Finished end-to-end transaction trace records **123** are forwarded by the transaction trace processor **102** to a transaction/topology inference processor **109**. The transaction/topology inference processor analyzes the received transaction record and queries known transaction related topology elements **112** from the reported topology model **105**. Those queried topology elements may include topology elements describing processes that executed portions of the transactions, host computing systems on which those processes were executed and communication activities of those processes. The inference processor **109** may, for transactions with outgoing calls for which no receiver side tracing data is available, like transaction T1 for which no trace data describing the

execution on application 2 is available, due to missing agents on host 2 **136** and application 2 **137**, use data describing the outgoing request to identify the topology element corresponding to the receiver component. In case of T1, the data describing the outgoing request identifies an incoming endpoint of the load balancing service **133** as receiving endpoint. The inference processor may use data describing the load balancing service as a component that does not process requests but only forwards requests, to determine that load balancing service **133** is not the receiver of the request but only an intermediate node.

(45) The inference processor may further use input/output mapping configuration data of the load balancing service **133** captured by the cloud and virtualization agent **152** to determine the network endpoint behind the load balancer **133** to which the incoming request from the execution of T1 **140** on application 1 **131** was forwarded. A load balancing service typically provides a set of input network endpoints and a set of output network endpoints. Incoming requests on input network endpoints are forwarded to output network endpoints according to the mapping configuration of the load balancing service.

(46) The analysis of the mapping configuration data of the load balancing service **133** in combination with the portion of the transaction trace data of transaction T1 **140** describing the outgoing request of application 1 **131** reveals that the final receiver of the request is a process (i.e. application 2 **137**) on host 2 **136**. The topology data provided by the cloud and virtualization agent **152** already showed the existence of host 2, and missing corresponding topology data from an OS agent showed that no OS agent and no transaction agents are installed on host 2.

(47) Only the inference of topology and transaction data showed that the not monitored host 2 is involved in the execution of transactions that are already, at least partial, monitored. A notification about a not fully monitored entity **114** may be issued to the user of the monitoring system as consequence of this finding.

(48) The inference processing of transaction T2 reveals a request from application 3 **135** to application 4 on host 4. OS agent **151c** is deployed to application 4, which provides data describing existence and type of application 4 on host 4 for the topology model. However, no transaction agent is deployed to the process running application 4, therefore no trace data describing the processing of T2 on application 4 is available. There are various reasons why no transaction agent is deployed to a process, reaching from the deliberate or accidental configuration that excludes a specific process from being monitored by a transaction agent to a process implementation and architecture for which no appropriate transaction agent is available. Software components providing database services typically use software architectures that do not support the injection of a transaction agent.

(49) The inference processor **109** may analyze the request sent from application 3 to application 4 to determine the network endpoint to which the request was sent. The OS agent **151c** deployed to host 4 **138** monitors processes that are executed on host 4 and their communication activities, including network endpoints which are used by those processes to receive requests. The inference processor **109** may use the topology data provided by OS agent **151c** to identify application 4 as the recipient of the request sent by application 3 to process transaction T2 **140**. The inference processor may further use topology data describing the process running application 4 to determine the type of this process. In case the process type of application 4 is one for which a transaction agent is available, it may issue a notification of an entity that is not fully monitored.

(50) Otherwise, it may use known topology and transaction facts to improve both transaction trace and topology data. As an example, the topology of application 4 may show that application 4 runs a database system of a specific vendor, type, and version, but it may not contain data identifying the network endpoints at which application 4 provides database services. The transaction trace data may show that application 3 sent a request for a database service to a specific network endpoint provided by application 4.

(51) The inference processor may improve the transaction trace data by adding data describing the

concrete database system to which the database service request was sent to the portion of the transaction trace data describing the database service request. It may further improve the topology model by adding data to the model that describes that application 4 provides a database service at the network endpoint to which application 3 sent its database service request.

(52) Improvements to the topology model identified by the inference processor **109** may be stored in a separate version of the topology model **106** that contains enrichments of the topology model identified by the inference processor. The reason for a separate, inference improved topology model is to avoid that the inference processor uses already inferred data for further inference processing as this improves the possibility of incorrect inference results.

(53) The enhanced end-to-end transaction trace data may be stored **111** in a transaction repository **103** for subsequent processing.

(54) An anomaly detection module **107** may use the enhanced end-to-end transaction traces **123** stored in the transaction repository **103** and the enriched topology model **106** to identify abnormal operating conditions of transaction executions and topology entities and to localize those abnormal operating conditions in the topology model to create topology localized anomaly data **124**. The created anomaly data may be used by a causality analysis module, together with topology data and other transaction trace data to calculate possible causal relationships between identified abnormal operating conditions to create networks of causally related anomalies **122**.

(55) Referring now to FIG. **1b** which shows a monitored environment containing various proxy components. A proxy component represents an intermediate communication component that terminates network connections on lower, transport-oriented network levels while simulating a continuous logical network connection on a higher, application level. A typical task of proxy component is to cache response data to reduce the load on components behind the proxy. In case a request is received by the proxy, it first checks if it received an equivalent request before and has the corresponding response in its cache. In case the response is cached, the cached version is returned to the requestor. Only if the response is not in the cache, the request is forwarded to the real receiver. Proxy components dedicated to the processing of HTTP based communication may adapt received requests before forwarding them with data indicating that the request was forwarded by a proxy component. Proxy components may use the X-Forwarded-For feature to store data identifying the network node that sent the request that was forwarded by the proxy (this may be another proxy component in case of a sequence of proxy components sequentially processing received requests. In this case, all proxies may use the X-Forwarded-For to store data describing the network host from which they received a request, and the sequence of data stored in the X-Forwarded-For data store reflects the sequence of the proxies forwarding the request).

Alternatively, proxies may use the “Forwarded HTTP Extension” feature to modify processed requests, by adding data describing the processing proxy and data describing the network node (i.e. IP address or domain name) or network endpoint (i.e. IP address or domain name and port) and an identifier for the protocol used to transfer the request.

(56) The functionality of a proxy component may be implemented in form of specialized hardware appliances, or it may be implemented as a software component.

(57) The exemplary monitored environment that is shown in FIG. **2** consists of an external host **157**, running outside of a data center or other environment controlled by the operator of the monitored environment, on which a web browser **158** is running. A browser agent **156** is injected into the web browser. The content displayed by the web browser may be provided by an application running in the monitored environment and the application may be configured to insert a browser agent **156** into provided content. Browser agents are inserted into the content they are monitoring. Therefore, they have the same level of access to browser and content internal data as process or transaction agents have to process internal data.

(58) The browser agent **156** monitors and traces **155** user interactions performed on the browser and sends monitoring and tracing data in form of topology localized trace data **120** to the

monitoring server. For browser side transaction monitoring data that describes portions of transaction executions that are performed outside the monitored environment, the topology localization may describe a localization aspect that is independent of the internal topology of the monitored environment. Topology localization aspects for browser-based portions of transaction executions may include data describing the geolocation of the web browser which was used to perform the transaction execution.

(59) An external proxy component **159** receives requests sent by the browser **158** on a network connection between the browser **158** and the external proxy component, adapts **168a** the received request (e.g. by using either the X-Forwarded-For or the “Forwarded HTTP Extension” feature as described above) to notify receivers of the request that the request was intermediately processed by the proxy and forwards the request to application 4 **161** running on host 4 **160** using a separate network connection between the external proxy component **159** and application 4 **161**.

(60) A dedicated proxy appliance **162** (i.e. a dedicated hardware component solely providing proxy functionality) receives requests sent from application 4 **161** on host 4 **160**, adapts **168b** those requests and sends them on a separate network connection to a proxy application **164** running on proxy host **163**. The proxy application **164** also adapts the received request and forwards it to application 5 **166** running on host 5 **165**.

(61) OS agents **151c**, **151d**, and **151e** are deployed to host 4 **160**, proxy host **163** and host 5 **165**. Those OS agents monitor processes running on the hosts they are deployed to and provide topology data **121** describing those hosts and the processes running on them. OS agent **151d** running on proxy host **163** monitors the proxy application process **164**, identifies the type of the process as proxy application e.g. by analyzing the command line used to start the process or by analyzing the executable used by the process. Further, the OS agent may analyze configuration data of the identified proxy application process to e.g. determine mapped input and output network endpoints of the proxy application process. This determined mapping data may be sent to the monitoring server **101** as part of topology data **121**. Common proxy applications are based on application architectures and/or programming languages that do not support the injection of in-process agents, therefore no transaction agents may be deployed to proxy application processes **164**.

(62) Transaction agents **150c** and **150e** are deployed to application 4 **161** and application 5 **166**. Those transaction agents monitor the execution of transactions by those applications and send the corresponding topology localized trace data **120** to the monitoring server.

(63) Both transaction agent **150c** and **150e** and browser agent **156** may add transaction tagging and correlation data to outgoing requests and may read transaction tagging data and correlation data from incoming requests. The read transaction tagging and correlation data may be sent to the monitoring server **101** as part of the topology localized transaction trace data **120** and may be used to identify and combine transaction trace data describing portions of the same transaction that were provided by different agents.

(64) Transaction T3 **167** is currently executed by browser **158** and the monitored environment. The execution of transaction T3 starts with a user interaction performed on the browser **158**, which is monitored by the browser agent. The browser agent reports the portion of T3 executed on the browser **158** in form of topology localized trace data to the monitoring server **101**.

(65) The transaction trace data **120** generated by the browser agent **156** may also contain data describing requests sent by the browser to process the transaction T3. This may include data describing the client and server-side network endpoints (i.e. network address and port) of the network connection used to send those requests.

(66) The request sent by the browser **158** to execute transaction T3 is received by the external proxy component **159**, adapted by the proxy component by adding data describing the network node from which the forwarded request was sent (i.e. by using the X-Forwarded-For feature and adding the data identifying the external host **157** as sender of the forwarded request) or in addition also adding data identifying the external proxy and the protocol used to transfer the request (i.e. by

using the “Forwarded HTTP Extension” feature).

(67) The external proxy component forwards the request to application 4 running on host 4. The transaction agent **150c** injected into application 4 recognizes the receipt of the request and creates transaction trace data **120** describing the receipt of the request which is sent to a monitoring server. The transaction agent **150c** captures data describing the client and server-side network endpoints used to transfer the received request (i.e. the network endpoint used by the external proxy as client-side endpoint to forward the request and the network endpoint used by application 4 to receive the request) and data describing the request adaptations performed by the external proxy component (i.e. by checking if the request contains X-Forwarded-For or “Forwarded HTTP Extension” data and capturing this data if it is available). Both network endpoint data and proxy request adaptation data are sent to the monitoring server as part of topology localized trace data **120**.

(68) The monitoring server may correlate portions of T3 **167** executed on the browser with portions of T3 executed on application 4, and may later perform topology inference processing with the data describing the communication between browser and application 4 to execute transaction T3. This inference may reveal that the communication between browser **158** and application 4 is not direct, but routed via at least one intermediate component. This may be performed by comparing the network endpoint data captured on the browser side with the network endpoint data captured by the transaction agent **150c** deployed to application 4 and determining that endpoint data captured on the server and client-side of the communication do not match. In addition, inference processing may check if data is available on request message data captured on the server-side of the communication that indicates that the intermediate components contain a proxy component (i.e. X-Forwarded-For or “Forwarded HTTP Extension” data). The inference process may then update the inference findings enriched topology model **106** with data of a proxy component **159** between the web browser **158** and application 4. A causality analysis module may use this enriched topology model to e.g. identify the external proxy component as a potential root cause of observed anomalies like a drop of the transaction load seen on application 4 combined with an increase of failing requests from browsers communicating with application 4 via the external proxy component **159**.

(69) Transaction T3 is partially processed by application 4 on host 4 and then continued with a request sent to application 5 running on host 5. Request from application 4 to application 5 to fulfill T3 are routed via the internal dedicated proxy appliance **162** and the proxy application **164** running on proxy host **163**.

(70) The transaction agent **150c** injected into application 4 **161** recognizes the request sent from application 4 to process T3 and creates transaction trace data describing this request from the client-side, including client and server-side communication endpoint data as seen from application 4 **161**. The request is sent to the dedicated proxy appliance **162** (which is seen as server-side communication endpoint from application 4), which adapts **168b** the received request (i.e. by adding X-Forwarded-For/Forwarded HTTP Extension data to the request) and forwards the request on a separate network connection to the proxy application **164**, which also adapts **168c** the received request and forwards it to application 5 running on host 5 using again another network connection. The transaction agent **150e** injected into application 5 recognizes the receipt of the request sent from application 4 to execute T3 and captures transaction trace data describing the receipt of the event. This included capturing data describing network endpoint data of sender and receiver side of the request as seen from application 5 (which sees proxy application **164** as sender of the request) and capturing data describing the received request from the server-side which may include data describing adaptations of the request performed by intermediate proxy components (X-Forwarded-For/Forwarded HTTP Extension data). The captured transaction trace data **120** is sent to the monitoring server **101**. The monitoring server correlates transaction trace data portions describing transaction T3 sent from transaction agents **150c** and **150e** and afterwards performs topology inference processing with the portion of the end-to-end transaction trace data describing the communication between application 4 **161** and application 5 **166**.

(71) Topology inference processing may again first analyze the communication endpoint data provided by the transaction agent **150c** deployed to application 4 **161** and by transaction agent **150e** deployed to application 5 **166**. This analysis reveals that the endpoints reported by the different agents do not match which indicates that the communication between application 4 and application 5 is not performed directly but routed via at least one intermediate node.

(72) Afterwards, inference processing may analyze the request message received by application 5 (which may have been captured by transaction agent **150e** and sent to the monitoring server **101** as part of topology localized transaction data **120**). The request message received by application 5 contains the adaptations performed by dedicated proxy appliance **162** and proxy application **164**, which may be used by the inference process to determine that the request sent from application 4 to application 5 was processed by at least two proxy components.

(73) For further considerations, it is assumed that request adaptations performed by proxy components (i.e. dedicated proxy appliances or proxy applications) use the X-Forwarded-For feature, as X-Forwarded-For is widely used in the industry and because X-Forwarded-For provides a subset of the data that is provided by the “Forwarded HTTP Extension” feature. As a consequence, all analyses and processing described here assuming the usage of the X-Forwarded-For also work for the “Forwarded HTTP Extension” feature.

(74) The analyzed request data may contain a X-Forwarded-For header field with a value starting with an identifier of the network node that originally sent the request (i.e. network address of host 4 **160** running application 4 **161**) which was added by the dedicated proxy appliance **162** to the request, followed by an identifier of the first proxy component forwarding the request (i.e. network address of dedicated proxy appliance **162**) which was added by the proxy application **164** running on proxy host **163**. There is no data contained in the X-Forwarded-For header identifying proxy application **164** as the second forwarder of the request as X-Forwarded-For only contains data identifying the entity from which a forwarded request was received. Therefore, the last element of a proxy chain is not represented in X-Forwarded-For data. To identify the last element in the proxy chain, inference processing may use the network endpoint data of the entity that sent the request to application 5, as provided by transaction agent **150e**.

(75) Network endpoint data typically contains a data identifying network node (i.e. a network address) and data identifying a component on the component on the network node (i.e. a port number identifying a socket on the network node) that is associated with a component (e.g. a process) on the network node that sends and receives data by the identified network endpoint.

(76) The network endpoint data of the request sender as seen by application 5 identifies proxy host **163** as network node from which the request was sent. The network endpoint data further contains a port number that identifies the network socket on proxy host **164** that was used to send the request. The operating system running on proxy host **163** manages the usage of sockets by processes (like the process running the proxy application **164**) and maintains mapping data between process identifiers (i.e. PIDs) and socketed identifiers (i.e. port numbers assigned to those sockets) of sockets used by the process. The OS agent **151d** deployed to the proxy host **163** may capture this mapping data and send it to the monitoring server as part of the topology data **121**. The OS agent may further analyze command line and executable binary used to start the proxy application process **164** to determine that this process provides a proxy service and it may further fetch configuration data used by the proxy service, like data defining the mapping of input and output network endpoints of the proxy service. This data may also be sent to the monitoring server as part of the topology data.

(77) The topology data processor **104** of the monitoring server **101** may process this data describing the proxy application process **164** to add an entity representing the proxy application to the reported topology model **105**. It may further store the proxy service configuration data fetched by the OS agent in the entity representing the proxy application, and the topology data may further store the data describing the usage of ports/sockets by the proxy application in the corresponding

topology entity.

(78) The inference processing may use the reported topology data describing proxy application **164** together with the sender side network endpoint data of the request received by application 5 to identify proxy application **164** as the last element of the proxy chain between application 4 and application 5. Inference processing may update the end-to-end transaction data describing transaction T3 to describe that communication between application 4 and application 5 to execute T3 is routed via proxy application **164**.

(79) In addition, inference processing may add data to the inferred topology model **106** describing that the service provided by application 5 that was used to process transaction T3 is connected to proxy application **164**. The inference processing may further analyze the request adaptations (i.e. X-Forwarded-For header) added by proxy application **163**, which identify the network node from which the proxy application **163** received the request. Analysis of the request adaptations provides the network address of the dedicated proxy appliance **162**. Inference processing may use the network address extracted from the request adaptation data to query the reported topology model for an entity with a matching network address.

(80) As no agent is (or can be) deployed to dedicated proxy appliance **162**, no topology data describing or indicating the existence of dedicated proxy appliance is available in the reported topology model. Inference processing may further compare the network address extracted from the request adaptations with the sender network address reported by the transaction agent **150c** deployed to application 4 to verify the hypothesis that the sender from which proxy application **164** received the request is already the original sender application 4 **161** on host 4 **160**.

(81) As this is not the case, this hypothesis is rejected and inference processing continues by analyzing the request adaptation data, which may reveal an entry added by dedicated proxy appliance **162** indicating that dedicated proxy appliance **162** received the request from host 4 **160** which executes application 4 **161**, the original sender of the request as reported by the end-to-end transaction trace data describing transaction T3 **167**. This supports the hypothesis of an unknown proxy component situated between application 4 and proxy application **164**.

(82) In addition, inference processing may compare the network address of the receiver network endpoint of the request as reported by transaction agent **150c** deployed to application 4 with the network address from which proxy application **164** received the request. Those network addresses are identical which also supports the hypothesis of an unknown proxy component situated between application 4 and proxy application **164**. As a consequence, inference processing may update the portion of the end-to-end transaction data of T3 describing the communication between application 4 and application 5 to indicate that this communication was routed via a proxy component with the address extracted from request adaptation data before it was routed via proxy application **164**.

(83) In addition, the inferred topology model **106** may be enriched with an entity describing dedicated proxy appliance **164**, and with data describing the observed communication activity of dedicated proxy appliance **164** with application 4 and proxy application **164**.

(84) It is noteworthy that request adaptation by proxy services is optional, and some proxy services may be configured to not perform it. As a consequence, further proxy service components may exist between application 4 **161** and dedicated proxy application **162**, between dedicated proxy application **162** and proxy application **164** or between proxy application **164** and application 5 **166**. Inference processing may not be able to identify those hidden proxy services, but it may use reported topology data in combination with reported request adaptation data and network endpoint data to determine if such hidden proxy service components exist.

(85) As an example, inference processing may reveal that the original sender of a request reported by request adaptation data is not identical with the original sender of the request reported by end-to-end transaction data. This indicates that at least one hidden proxy service component is located between the original sender of the request as reported by the end-to-end transaction trace data and the first proxy service component that performed request adaptation.

(86) OS agents monitor communication activities of processes executed on operating systems on a TCP/IP level. This also includes communication activities that are not captured by transaction agents. Therefore, the absence of monitoring data indicating TCP/IP level communication between two processes or hosts that were reported as adjacent entries in a proxy chain by request adaptations may be used as an indication that one or more hidden proxy service components, that did not perform request adaptations, are located between the two reported proxy components.

(87) Other contradictions between transaction agent provided transaction-based communication data, like reported client-side (i.e. sender of a request) and server-side (i.e. receiver of a request) sender and receiver network endpoint data and captured proxy request adaptation data (i.e. X-Forwarded-For header data), and OS agent provided network-based (i.e. TCP/IP connection data), may be used as indication of hidden proxy service components (i.e. proxy service components that do not perform request adaptations). As an example, request adaptation data may indicate two adjacent proxy applications in a proxy chain, both proxy applications may run on different host computer systems, monitored by OS agents. The OS agents monitor and report the network communication activities of processes running on both hosts, including the network communication of the proxy application processes. Missing direct network communication data between the proxy applications may be used as an indicator that communication is routed via at least one hidden proxy service component between the two OS agent monitored proxy applications, despite captured request adaptation (X-Forwarded-For) data indicates that both proxy applications communicate directly.

(88) Coming now to FIG. 1c which depicts a variant of a monitored environment in which container-based virtualization and process isolation (as e.g. provided by the docker product is used. Containers are software components executed by a host operating system and provide an isolated operating environment for processes running inside such containers. Containers simulate a separate operating system using resources of the host operating system like network interfaces, CPUs, main and secondary storage. A container configuration defines a container specific view and usage permissions of resources of the host operating system.

(89) The network related portion of the container configuration may be used to define network addresses for containers (i.e. a container running on a host computer system with network address “a” may be assigned a network address “b” and the container management system assures that the container is reachable via address “b”), define port mappings between container internal and external ports (i.e. a port **80**, used by a process running in a container may be mapped to an external port **8080** of the host operating system, and the container management system assures that communication from inside the container addressing port **80** is forwarded to port **8080** of the host operating system and vice versa), and to configure container specific overlay networks that define connectivity, visibility and accessibility amongst different container instances.

(90) OS agents (i.e. OS agents **151f** and **151h**) deployed to hosts on which container management systems are running (i.e. host 6 **181** and host 8 **177**) may also monitor running containers (i.e. container 1, **169**, container 2 **171**, container 3 **173** and container 4 **178**), running container manager processes (not shown) and running container manager coordination processes (to e.g. ease inter-container communication between containers running on different host computer systems).

(91) OS agents may also report container related topology data identifying and describing running containers, where the OS agent may further fetch above described container configuration data and add it to the topology data which is sent to the monitoring server. OS agents may further fetch configuration data potentially affecting multiple containers, like configuration data defining container overlay networks and send this data to the monitoring server as part of topology data.

(92) The topology data processor **104** of the monitoring server **101** receives this container related topology data and enriches the reported topology model with topology entities describing the reported containers, their configuration data and network mapping (network address, port) and network connectivity, visibility and accessibility data. The created topology model describes which



containers are hosted by which host computer systems and how those containers are accessible from other topology elements.

(93) The OS agents may further manipulate the configuration of starting containers in a way that transaction agents deployed to processes running in the containers have access to the monitoring server (either directly or via the OS agent running on the host), and in a way that transaction agents running in the container have access to identification data of the host computer system and the container, e.g. by creating a file containing this identification data on a position and with a name in the file system of the container, where position and name are known by transaction agents deployed to processes running in the container. Transaction agents running in container environments may use this identification data for the creation of topology localized transaction trace data to create data describing the location of monitored transaction executions in terms of a container instance in which the execution occurred and in terms of a host computer system running the container instance.

(94) The scenario of FIG. 1c shows a monitoring environment containing the hosts host 6 **181**, host 7 **175** and host 8 **177**, all those hosts being monitored by OS agents **151f**, **151g**, and **151h**.

(95) Host 6 **181** runs a container management system which currently executes containers 1 **169**, container 2 **171**, and container 3 **173**. OS agent **151f** monitors **153** the running containers and the processes running in those containers (i.e. application 6 **170** in container 1 **169**, application 7 **172** in container 2 **171** and application 8 **174** in container 3 **173**) and sends corresponding topology **121** data to the monitoring server **101**. Transaction agents **150f** and **150g** are injected into processes running in those containers which monitor transaction executions performed on those processes and send the corresponding topology localized transaction trace data **120** to the monitoring server **101**.

(96) OS agent **151g** is running on Host 7 **175**, which monitors **153** processes executed on host 7 and sends corresponding topology data **121** to the monitoring server **101**. A transaction agent **150h** is deployed to process application 9 **176**, to monitor transaction executions.

(97) Container 4 **178** is currently running on Host 8 **177** and application 10 **79** is executed in container 4. An OS agent **151h** monitors **153** the container and the application process running inside the container and sends corresponding topology data **121** to the monitoring server.

(98) Each container has its specific view of the underlying computer network, as defined in the container configurations of the individual containers, which is depicted in FIG. 1b as container network isolation **180a** to **180d**. Data describing those container specific views of the underlying network are also part of the topology data and the container specific network views are also described in the reported topology model **105**.

(99) Transaction T4 **182** is executed and enters the monitored environment via the network isolation **180a** of container 1 **170**. The sender of the incoming request to execute a portion of T4 may use an externally visible port to communicate with application 6 in container 1, which may be translated into a container internal port by the container management system. The translation of external ports to container internal ports may be performed according to a port mapping configuration that specifies a set of container internal ports that are mapped to container external ports. This port mapping data may be defined for each container and stored in container configuration data. An OS agent may read this port mapping data and send it to a monitoring server **101** as part of topology data **121**. The topology model **105** may contain port mapping data for each container represented in the topology model. Application 6 receives the incoming request at the internal port, and the transaction agent **150f** deployed to application 6 recognizes an incoming request from the container internal port. The agent **150f** may monitor the execution of transaction 4 on application 6 and send the corresponding topology localized trace data to the monitoring server.

(100) Inference processing on the monitoring server may analyze the tracing data describing the incoming request on application 6, use corresponding topology localization data to determine that application 6 runs in container 1, and use the network configuration data stored in the topology to find the corresponding external port for the container internal port reported by agent **150f** as port

used by the sender of the request to process transaction T4. The external port may together with the network address be used to identify the topology entry representing the component (i.e. a specific process) that sent the request that initialized the execution of T4 on application 6. T4 may spawn on application 4 into branch **182a** and **182b**. Branch **182a** may continue with a request from application 6 to application 7 running in container 2. The request is first routed via the network isolation of container 1, which e.g. maps a container internal port into an external port, and then routed via the network isolation of container 2, which may map the external port to a container internal port of container 2. The agent **150f** deployed to application 6 recognizes an outgoing call to a port internal to container 1 and the agent **150g** deployed to application 7 recognizes an incoming call to a port internal to container 2. Corresponding topology localized transaction trace data may be sent to the monitoring server, and the inference processing on the monitoring server may use container network configuration data stored in the topology to reconstruct the path of the communication between application 6 and application 7 via the container network isolation of container 1 and container 2.

(101) Application 7 may, to further process branch **182a** of transaction T4 send a request to application 8 **174** running in container 3 **173**. As the communication between application 6 and application 7, also this request is routed via the network isolation of both containers, causing two port mappings. However, in this case, no transaction agent is deployed to application 8. Agent **150g** may send topology localized transaction trace data to the monitoring server and inference processing may use network configuration data of container 2 and container 3 to determine the internal port of container 3 addressed by the request. Inference processing may further analyze the type of the request sent by application 7 to application 8 to identify the type of the requested service (i.e. database service) and update the topology entity representing application 8 in the inferred topology model to provide the requested service at the previously identified internal port of container 3. Inference processing may further use other topology data describing application 8 to update the portion of end-to-end transaction traced data of T3 **167** that describes the communication between application 7 and application 8 (e.g. by adding data indicating the type and vendor of the database application).

(102) The branch **182b** of T4 is continued by a request sent from application 6 **170** to application 9 **176** running on host 7 **175**. As the sender of the request (i.e. application 6) is running in a container environment, and the receiver of the request (i.e. application 9) runs outside the container environment, the sender endpoint data of the request identifies a network endpoint that may not be visible for the receiver. The container management system and the networking system connection host 6 and host 7 assure that the request and the corresponding response are transferred between application 6 and application 9, even if the sender application is not directly visible for the receiver application.

(103) Agents are deployed to application 6 and application 9, which monitor the execution of branch **182b** of transaction T4. The agent **150f** deployed to application 6 recognizes the sending of a request to a network endpoint (i.e. the address of host 7 and a specific port), and the agent **150h** deployed to application 9 recognizes the receipt of a request from a network endpoint that is not visible from the network context of application 9.

(104) Transaction agent **150f** may add correlation data to the request, which is read by agent **150h**, and both agents may add this correlation data to transaction monitoring data, e.g., in form of incoming and outgoing trace records **470** and **480** as described in FIG. 4. The trace records are sent to the monitoring server **101** and the transaction trace processor **102** may use the correlation data to identify corresponding incoming and outgoing trace records to create end-to-end transaction trace data. The network endpoint data contained in the outgoing trace record **470** sent from agent **150f**, to report the branch **182a** of transaction T4, is consistent with the network endpoint data contained in the network endpoint data contained in the incoming trace record **480** sent by agent **150g**, as both sender and receiver agents are situated in container environments managed by the same

container management component and may therefore share the same network isolation component. Therefore, network addresses and ports seen by application 6 and application 7 may be consistent. It should be noted that network isolation configurations are possible which define different network environments and different network addresses and ports for containers managed by the same container management component.

(105) However, the reported endpoint data, sent by agent **150f** to report the branch **182a** of transaction T4, is not consistent with the network environment seen by host 7 and application 9, as the sender endpoint is not directly visible because it is located in container 1.

(106) The transaction/topology inference processor **109** may analyze data stored in the reported topology model **105** showing that application 6 runs inside container 1 in combination with transaction trace data describing the communication of application 6 and application 9. This analysis may enrich the transaction trace data portion describing the sending of the request from application 6 to application 9 with data describing that the sender application 6 is executed inside container 1 and the receiver application 9 is executed outside the container environment.

(107) To continue with branch **182b** of transaction 4, application 9 **176** may send a request to application 10 **179** running in container 4 **178** on host 8 **177**. Container 4 **178** is isolated from the network environment visible to application 9 by its container network isolation **180d**. The container network isolation **180d** may be configured to map a container internal port used by application 10 **179** to provide the service requested by application 9 **176** to a container external port. Application 9 may send the request to the mapped container external port on host 8 and the container network isolation **180d** may forward the request to the internal port used by application 10. Application 10 processes the request and sends it back to application 9 via the container network isolation **180d**.

(108) The agent **150h** deployed to application 9 recognizes an outgoing request directed to a port (i.e. external port mapped to the container internal port used by application 10 to provide the service requested by application 9) on host 8 and sends corresponding transaction trace data to the monitoring server **101**.

(109) An OS agent **151h** is executed on host 8 which monitors all running processes and all containers running on host 8 and all processes running in those containers and sends corresponding topology data to the monitoring server. In addition, OS agent **151h** reads container configuration data of running containers including configuration data of the container network isolation, and data describing communication activities of monitored processes, like the ports on which processes are listening for incoming requests.

(110) No transaction monitoring agent is deployed to application 10, e.g. because application 10 is implemented using a technology for which currently no matching agent is available. An example for application types for which no matching agent may be available are applications providing database services.

(111) The transaction/topology inference processor **109** analyses transaction trace data describing sending of the request sent from application 9 to application 10 and recognizes that no corresponding transaction trace data describing the receipt of the request is available. As a next step, transaction/topology inference processing may use receiver network endpoint data reported by the agent **150h** deployed to application 9 which sent the request. The reported receiver endpoint data contains a network address that identifies host 8 as computer system receiving the request. The reported endpoint data also contains a port that identifies the process to which the request is directed. An analysis of the topology data describing host 8 and all processes running on host 8 shows that no process is running directly on host 8 that listens on the port of the reported receiver network endpoint data. A further analysis of the reported container network isolation configuration data may show that the reported port is mapped to a container internal port on container 4 **178**, and a further analysis of the processes running in container 4 and the container internal ports used by those processes may show that application 10 **179** listens on the mapped container internal port, which identifies application 10 **170** in container 4 **178** as receiver of the request sent by application

9.

(112) As a consequence of identifying the process that received the request sent from application 9 as application 10, transaction/topology inference processing may enrich the portion of the transaction trace data of transaction T4 that describes the sending of the request from application 9 to application 10 with corresponding topology data.

(113) As an example, data describing the outgoing request may identify the request as a request for a database service, but the monitoring data may not contain data describing details of the requested database service, like vendor, type or version of the database management system providing the database service. OS agents may, as part of process monitoring activities, read data describing running processes, like the command line used to start processes, installation and working directories of processes and configuration data used by those processes. This process description data, which is sent to the monitoring server as part of topology data, may also describe vendor, type, and version of a database management system or another type of service provided by monitored processes. The topology data processor may enrich the topology model with this process description data and transaction/topology inference processing may use this process description data of processes involved in the execution of monitored transactions to enrich the transaction trace data.

(114) Transaction/topology inference processing may in addition enrich the topology model with data derived from the processed transaction trace data, e.g. by adding data to the portion of the topology model describing application 10 **179** and container 4 that specifies that application 10 provides a database service at the previously identified container internal port.

(115) Referring now to FIG. 2, which depicts flowcharts of processes performed by OS agents and cloud/virtualization agents to generate topology data and report it to a monitoring server. FIG. 2a describes the process and process communication activity monitoring as performed by an OS agent, FIG. 2b describes the processing of identified changes of a cloud or virtualization environment by a cloud/virtualization agent and FIG. 2c describes the processing of a detected container startup by an OS agent.

(116) The processing performed by the OS agent to monitor processes running on an operating system and the communication activities of those processes starts with step **200** when an operating system with an installed OS agent starts. The OS agent is typically configured to automatically start with the start of the operating system. Following step **201** fetches or creates a unique identifier for the monitored operating system and fetches network address data from the host. The unique topology identifier may be deferred from network address data of the host and other data like e.g. a timestamp of the installation of the OS agent on a monitored host computer system to generate a value that is unique for each monitored host computing system. The OS agent may, on startup, check a known persistent storage location if such a stored identifier exists, and in case it exists, uses it. Otherwise, it may create the identifier, store it in the known persistent storage for the next start and use the created identifier.

(117) Subsequent step **202** starts monitoring the startup and shutdown of processes running on the operating system. Monitoring process startup and shutdown may be performed by cyclically polling management data structures of the operating system that list the currently running processes. This data may be used to create topology data describing the processes executed by the operating system of the monitored host computing system.

(118) Following step **203** may manipulate the process startup procedure used by the operating system to check on process startup if a transaction agent matching the type of the starting process is available and in response to an available transaction agent, inject the transaction agent into the starting process.

(119) Afterwards, step **204** may start to continuously monitor communication activities of processes running on the monitored system. Monitoring of communication activities may be performed by cyclically polling operating system data structures that list the currently used network

endpoints (i.e. ports) per running process. Some variant embodiments may only monitor long-living server-side listen ports (those ports are typically used by server processes providing services on a specific port and are typically reserved during the whole lifetime of the process, regardless if network communication using the port is ongoing). Other variants may also monitor short living client-side ports (those ports are only reserved during currently ongoing network communication). Yet other embodiments may instrument the network endpoint/port management component of the operating system to get notified on activities like reservation, freeing of ports by processes, setup or shutdown of network connection by processes or the data transfer over network connections, to provide more precise and detailed monitoring data of the network activities of the monitored processes. In addition to active communication activity, step **204** may also monitor the establishment and shutdown of server-side communication endpoints (i.e. listen ports on which server processes receive requests).

(120) Following step **205** starts to cyclically send OS topology messages containing process execution and communication-related topology data of processes running on the monitored operating system. The process then ends with step **206**.

(121) The processing of received cloud/virtualization environment change events by a cloud/virtualization agent **152** is shown in FIG. **2b**. Virtualization and cloud systems typically provide notification services that report changes of the cloud/virtualization environment. Those notifications may contain but are not limited to notifications reporting startup or shutdown of a cloud/virtualization node, changes of the configurations of a cloud/virtualization node and changes of the network configuration of the cloud/virtualization environment. Cloud/virtualization agents **152** may be registered as listeners for such notifications and may process such notification events as conceptually described in FIG. **2b**.

(122) The process starts with step **210** when a new environment change notification is received by a cloud/virtualization agent. Following decision step **211** determines whether the notification reports the start of a virtual host in the monitored cloud/virtualization environment. In case the start of a virtual host is notified, step **212** is executed, which may use data provided by the notification to determine a unique topology identifier for the host, fetch network address data of the started host and other data describing the host computer system, like type, specification and number of CPUs, size of main and secondary memory of the host computer system and type and version of the operating system running on the starting host computer system. In addition, the cloud/virtualization agent may fetch network isolation data, describing the network environment of the starting host computer system, like data describing a virtual private cloud (VPC) environment containing the started virtual computer system. A VPC typically isolates a portion of a cloud environment by defining access and visibility rules for incoming and outgoing communication between components inside and outside the VPN. The cloud/virtualization agent may either extract this data from the received notification, or it may, on receipt of a notification describing the start of a virtual host, access a cloud/virtualization management unit to request this data.

(123) Following step **213** creates a virtual host topology message **430** using the data gathered in step **212** and sends it to a monitoring server **101**.

(124) Step **214** determines whether the received notification indicates the start of a virtual network traffic dispatching service.

(125) A virtual network traffic dispatcher, like an Amazon AWS Elastic Load Balancer (ELB), typically receives incoming traffic on one or more network endpoints (i.e. network address and port) and forwards incoming traffic to one or more receiving network endpoints. Conceptually, the functionality of a virtual network dispatcher is defined by one or more network mapping tables that define how incoming traffic is distributed amongst receiving endpoints. Such a mapping table typically contain one incoming network endpoint and multiple outgoing network endpoint that receive incoming data. The virtual traffic dispatcher may distribute incoming traffic evenly to mapped receiving network endpoints, and it may, in addition, use specific dispatching rules to route

incoming traffic. Such dispatching rules may connect specific parameters of incoming requests with certain receiving endpoints, and the virtual network traffic dispatcher may, on receipt of an incoming request, analyze the request and apply those dispatching rules to determine a receiving endpoint of the incoming traffic. Virtual network dispatchers are typically located at the boundary of cloud/virtualization environments to receive incoming traffic from the outside world and dispatch it to components of the cloud/virtualization environment in a controlled way.

(126) In case the received notification reports the start of a virtual network dispatcher, step **215** is executed, which fetches data identifying and describing the virtual network dispatcher. The fetched data may include but is not limited to an identifier uniquely identifying the virtual network traffic dispatcher, data describing the mapping of incoming network endpoints to outgoing network endpoints, data describing traffic routing rules and data describing the network environment of the virtual network dispatcher, like data describing a VPC to which the dispatcher belongs.

(127) Following step **216** creates a traffic dispatcher topology message **440** using the data extracted in step **215** and sends it to a monitoring server **101**.

(128) Subsequent step **217** checks if the received notification indicates the start of a virtual endpoint service in the cloud/virtualization environment. Cloud/virtualization environments may provide specialized network nodes that are dedicated to the provision of specific services, like file storage or database services. Those network nodes typically provide no access to an underlying host computer system. Therefore, it is typically also not possible to deploy OS agents or transaction agents to those virtual endpoint services. In case step **217** determines that the received notification reports the start of a virtual endpoint service, step **218** is executed which gathers data identifying and describing the started virtual endpoint service, which is followed by step **219** which creates a virtual endpoint service topology message **450** using the data fetched by step **217** and sends the created virtual endpoint service topology message to a monitoring server **101**. The process then ends with step **220**.

(129) It is noteworthy that cloud/virtualization management and monitoring entities may send additional types of notifications, reporting other changes of the cloud/virtualization environment. The cloud/virtualization agent may handle those additional notification types in a similar way as the notifications shown in FIG. **2b**, by determining an identifier for the affected cloud/virtualization entity, fetching data describing the affected cloud/virtualization entity and sending a corresponding topology message to a monitoring server which may then update its topology model accordingly. Such additional notifications may include but are not limited to notifications reporting the shutdown of a cloud/virtualization entity, notifications reporting a configuration change of a cloud/virtualization entity or notifications reporting the creation, modification or deletion of a network isolation area like a virtual private cloud.

(130) The process of monitoring and reporting starting container environments by the OS agent is shown in FIG. **2c**. The process starts with step **230** when the OS agent detects the start of a new container environment and continues with step **231** which modifies the configuration of the starting container to support the monitoring of container internal processes and the injection of transaction agents into container internal processes. Following step **232** determines a unique identifier for the starting container and fetches configuration data of the starting container, including data describing the network isolation configuration of the starting container. Subsequent step **233** creates a container topology message **420** using the data fetched by step **232** and sends the create topology message to a monitoring server **101**. The process then ends with step **234**. The container monitoring mechanism of the OS agent may also be configured to detect the shutdown of containers and send corresponding topology messages to a monitoring server, which may, on receipt of a message indicating the shutdown of a container, update a topology model by e.g. removing the entity describing the terminated container from the topology model.

(131) Referring now to FIG. **3**, which describes processes related to the monitoring and reporting of transaction execution activities by transaction agents **150a-150h**. The monitoring of process

outgoing transaction communication is described in FIG. 3a, the monitoring of process internal transaction processing is described in FIG. 3b and the monitoring of received incoming transaction communications by a process is shown in FIG. 3c.

(132) The monitoring and reporting of outgoing transaction communication activities starts with step **300** when an agent deployed to a process that executes a monitored transaction detects an outgoing communication performed by the monitored transaction. Outgoing communication activities may be detected by sensors placed in code portions that are executed to perform outgoing communication activities, like sending a request. Those sensors are executed whenever outgoing communication activities are performed and they report the occurrence of communication activities, including details of the communication activities to the agent **150**. In addition, those sensors may alter messages sent by the communication activities, e.g. by adding transaction correlation data to those messages. Following step **301** fetches local transaction correlation data which identifies the portion of the monitored transaction executed by the current process and subsequent step **302** captures the network endpoint data of sender and receiver of the outgoing communication (i.e. network address and port used by the sender to send the message and network address and port to which the sender process sends the message). Afterwards, step **303** determines the type of the outgoing communication activity and captures the communication type specific communication message. The type of the communication activity may be derived from the sensor that reported the communication activity, as different sensors may be deployed to code related to different types of transaction communication. The determined type of the communication activity typically enables the sensor or the agent to fetch and interpret the data send by the communication activity and to extract communication type specific message data.

(133) Step **304** captures data identifying and locating the process performing the outgoing communication activity within the topology model **105** maintained by the monitoring server **101**. This identifying data may contain but is not limited to data identifying the host computer system running the process, like the network address of the computer system and data identifying the process within the scope of the host computer system, like the command line used to start the process together with the name of the executable file used to execute the process, a process identifier (PID) of the process, a start timestamp of the process or a combination of some or all of them.

(134) Following step **305** creates a transaction trace message describing the outgoing communication activity, like an outgoing trace record **470**, initializes it using data fetched by previous steps and sends the created outgoing trace record **470** to a monitoring server **101** for analysis. The process then ends with step **306**.

(135) The process of monitoring process internal transaction execution activity as described in FIG. 3b start with step **310**, when the agent detects process internal transaction activity. Process internal transaction activities may be detected by sensors placed to code that performs no communication activity. Following step **311** fetches or creates correlation data identifying the process local execution of the monitored transaction and subsequent step **312** captures data describing the process local transaction activity, like e.g. the execution of a method that is instrumented with sensors. Step **313** afterwards captures data identifying and locating the process performing the monitored transaction execution within the topology model **105**. Further, step **314** creates a transaction trace message like a local trace record **460**, initializes it with data fetched by the previous steps and sends the local trace record **460** to a monitoring server for analysis. The process then ends with step **315**.

(136) The process of monitoring incoming communication activity, as described in FIG. 3c starts with step **320** when an agent deployed to a monitored process recognizes incoming transaction communication activity. Incoming transaction communication activity may be recognized by sensors instrumented to code that processes incoming transaction activity. Following step **321** fetches transaction correlation data from the incoming communication message, which may have

been attached to the communication message by step **301** of process **3a** and subsequent step **322** captures local visible network endpoint data of sender and receiver of the received incoming communication message (i.e. network address and port of sender and receiver of the incoming message as seen from the receiving process). Following step **323** creates local transaction correlation data identifying the process local transaction execution initiated by the received incoming transaction communication and step **324** determines the type of the incoming communication message (i.e. by using the specific sensor reporting the incoming communication message that is instrumented to code specific to a certain communication type) and captures the received communication type specific communication message. It is noteworthy that for a communication performed between two processes, the communication message sent by a sending process and e.g. captured by an agent injected into the sending process (i.e. step **303** of process **3a**) and the communication message received by the receiving process and e.g. captured by an agent injected into the receiving process (i.e. step **324** of process **3c**) may differ, because potential intermediate nodes that route and dispatch the message from the sender process to the receiver process may add data to the message. An example of such added data is the “X-Forwarded-For” header field that intermediate proxy entities may add to HTTP based communication messages. (137) Subsequent step **325** fetches data identifying and locating the process receiving the incoming communication message within the topology model and following step **326** creates a transaction trace message describing the incoming communication message, like an incoming trace record **480**, initializes the incoming trace record with data fetched by the previous steps and sends it to a monitoring server **101** for analysis. The process then ends with step **327**.

(138) Coming now to FIG. **4** which depicts data records that may be used to transfer topology and transaction trace data from agents to monitoring servers. An OS topology record **400** as shown in FIG. **4a** may be used to transfer data describing a concrete or virtual host computer system running a specific operating system and may contain and but is not limited to a topology identifier **401** that uniquely identifies the described host computing system within the topology model **105**, a network address data **402** field, containing one or more network addresses of the described host computer system and other payload data **403** further describing the host computer system e.g. by specifying the number, version and speed of CPUs of the host computing system, size of main memory, data describing secondary storage of the host computing system like hard disks, data describing the capacity of network connections and type and version of the operating system of the host computing system.

(139) A process topology record **410** as shown in FIG. **4b** may be used to transfer topology data describing a process running on a host computer system and may contain but is not limited to a topology identifier **411** identifying the process within the scope of the entity executing the process (i.e. virtual or concrete host computing system or container), a topology identifier **412** identifying the entity executing the process, a network activity and endpoint data section **413** containing data describing client and server-side communication endpoints operated by the process in form of network address, port and client/server indicator and additional payload data **414** further describing the monitored process, like type of the process in form of application an platform used by the process (i.e. Oracle Java® or Microsoft .NET® application platform), type of an application server environment executed by the process and resources assigned to the process, like CPU and disk usage quotas or maximum amount of main memory assigned to the process.

(140) It is noteworthy that a process may perform multiple network communication activities, with various counterparts, using different ports and addresses and with varying client/server role. Therefore, the network activity and endpoint data **413** may contain multiple entries, each entry describing a network communication of the process in form of a client/server role indicator of the process, and network addresses and ports of both network endpoints involved in the network communication. In addition, the network activity and endpoint data may contain data describing individual network communication activities in detail, consisting of data describing start and end



time of a network communication, amount of data transferred during the communication activity in both directions and data describing the quality of the network communication e.g. in terms of delay times or number of data packet resents due to network errors that occurred during the communication activity.

(141) FIG. 4c shows a container topology record **420** which may be used to transfer data describing a container environment executed on a host computer system. A container topology record **420** may contain but is not limited to a topology identifier **421** identifying the described container within the context of the entity containing it, like the host computing system executing the container, a topology identifier **422** of the entity containing and executing the container environment, a network mapping data field **423** containing data describing the network isolation configuration of the container, like mapped ports, container internal network address of the container, in case the container is connected to a container overlay network, name of the overlay network and data describing the overlay network. A container topology record **420** may also contain a payload data field **424** that further describes the container and that may contain data describing the type of the container environment or resources assigned to the container.

(142) A virtual host topology record **430** as shown in FIG. 4d may be used to transfer data describing a virtualized host computing system. A virtual host topology record **430** may contain but is not limited to a topology identifier **431**, identifying the described virtualized host computing system within a topology model **105**, a network address data field **432** containing at least one network address of the described virtualized host computer system, a field **433** containing additional OS topology entity matching data, that may be used to identify a corresponding topology entity reported by an OS topology record **400** sent by an OS agent that runs on the virtualized host computing system, and a payload data field **434** further describing the virtualized host computing system, by specifying number, version and speed of CPUs of the virtualized host computing system, amount of main memory, data describing available secondary storage capacity like number and capacity of hard disks and data describing type and version of the operating system running on the described virtual computer system.

(143) A typical monitoring setup of a virtualized or cloud environment contains a cloud/virtualization agent **152** which e.g. monitors the startup of virtual host computing systems and reports topology data describing those virtual host computing systems from a virtualization perspective, e.g. in form of virtual host topology records **430**, to a monitoring server, which enriches a topology model **105** of the monitored environment with data describing the reported virtual host computer systems. In addition, OS agents **151** may be deployed to the virtual host computing systems, which also provide topology data to a monitoring server, e.g. in form of OS topology records **400**, that describe the virtual host computing systems from an operating system perspective. The topology data processor **104** of a monitoring server **101** receives OS topology records **400** and virtual host topology records **430** and may use network address data **402/432** and optional additional OS topology entity matching data **433** to identify matching OS topology record/virtual host topology record pairs describing the same virtual host computing system. The topology data processor may merge such identified topology record pairs and use them to create one topology entity representing the virtual host computing entity that describes both its virtualization and OS aspects.

(144) Topology data describing virtual traffic dispatcher entities, like cloud load balancers may be transferred using virtual network traffic dispatcher records **440** as shown in FIG. 4e. A virtual network traffic dispatcher record **440** may contain but is not limited to a topology identifier **441** identifying the described virtual traffic dispatcher within a topology model, a network mapping data field **442**, e.g. in form of one or more mappings that map input network addresses and ports to one or more output addresses and ports and a payload data field **443** further describing the virtual traffic dispatcher, e.g. by its type or capacity (e.g. max number of requests per time interval).

(145) Virtual endpoint service records **450**, as shown in FIG. 4f may be used to transfer data that

describes virtual endpoint services, like database or file storage services provided by cloud/virtualization environments. A virtual endpoint service record **450** may contain but is not limited to topology identifier **451** identifying the described virtual endpoint service in a topology model, a network endpoint data field **452** containing data describing the network endpoints at which the provided services can be accessed and a payload data field **453** that further describes the virtual endpoint service by specifying its type (e.g. database or file storage service), or its maximum and currently used capacity.

(146) Data records that may be used to transfer transaction trace data are shown in FIGS. **4g**, **4h**, and **4i**. Transaction traced data describing process local transaction activities may be transferred using a local trace record **460** as shown in FIG. **460**. A local trace record **460** may contain but is not limited to a local correlation data field **461** uniquely identifying the portion of a monitored transaction performed by a specific process, a process topology identifier field **462** identifying the process on which the transaction execution was performed within a topology model **105** and a payload data field **463** containing data that further describes the monitored portion of the transaction execution, like names of called methods and execution performance data of those methods, resources used for the transaction execution, like allocated memory or required CPU cycles.

(147) Transaction execution activities that perform process outgoing communication activities, like the sending of requests to other processes, may be transferred using outgoing trace records **470** as described in FIG. **4h**.

(148) An outgoing trace record **470** may contain but is not limited to local correlation data **471** identifying the portion of the monitored transaction execution that performed the outgoing communication activity, a topology identifier **472** identifying the process executing the transaction, a client-side network endpoint data field **473** containing data describing the sender and receiver network endpoint (network addresses and ports) as seen by the sending process and a communication data field **474** containing data describing the type of the performed outgoing communication activity (e.g. by specifying the application-level protocol used to send the message, like HTTP) and the sent communication message (e.g., for HTTP based communication, the data describing the sent HTTP request).

(149) Data describing incoming communication activities, like requests received by processes, that start the execution of a portion of a monitored transaction by the received process, may be transferred using incoming trace records **480** as shown in FIG. **4i**. An incoming trace records **480** may contain but is not limited to a local correlation data field **481** uniquely identifying the portion of the transaction execution performed by the receiving process to handle the incoming communication activity, a sender correlation data field **482** containing an identifier for the portion of the monitored transaction performed by the process that sent the incoming communication message, a topology identifier **483** identifying the receiving process within a topology model **105**, a field **484** to store receiver side network endpoint data in form of network addresses and ports of sender and receiver of the communication activity as seen by the receiver process and a field **485** to store communication type and communication message as received by the receiver process.

(150) The topology and trace records described in FIG. **4** may be used to send topology and transaction monitoring data from agents deployed to the monitored environment (e.g. transaction agents **150**, OS agents **151** or cloud/virtualization agents **154**) to a monitoring server **101**.

(151) Referring now to FIG. **5** which shows data records that may be used by a monitoring server **101** to describe and store topology models and end-to-end transactions. FIGS. **5a**, **5b**, **5d**, and **5e** describe such records in detail. FIGS. **5c** and **5f** show their usage by an example.

(152) A topology node **500** as shown in FIG. **5a** may be used to represent a topological entity in a topology model **105** maintained by a monitoring server **101**. A topology node **500** may contain but is not limited to a topology identifier **501** identifying the topology node within the topology model, a parent topology node identifying **502** a parent or enclosing topology node of a given topology

node (e.g. for a topology node describing a process the parent topology node identifier would identify the topology node representing the host computing system on which the process is executed), an entity type field **503** specifying the type of the described topology entity (e.g. process, host computing system, virtual traffic dispatcher etc.), a type specific network data field **504**, containing network data that is available for the specific entity type (e.g. network addresses for host computing systems, network addresses and port for processes) and type specific additional data field **505** which contains data that further describes the topology entity (e.g. assigned resources, and data describing the functionality of the process, like executable name, command line to start the process and configuration data for process entities, network configuration data for container entities, available resources and configuration data for host computing system entities and network traffic mapping configuration data for network traffic dispatcher entities).

(153) Topology communication nodes **510** as shown in FIG. *5b* may be used to describe communication activities between topology entities. A topology communication node **510** may contain but is not limited to a client topology identifier **511** identifying the topology entity that initiated the communication, a server topology entity identifying the topology entity that received and served the communication activity and a communication type field **513** containing data further describing the communication activity (e.g. TCP/UDP in case only low-level communication data is available, or in case application level data is available from transaction monitoring, the application level protocol, like HTTP for web browser or web service related communication or JDBC/ODBC for database service related communication).

(154) A processing trace node **520** as shown in FIG. *5d* may be used to describe portions of a monitored transaction execution that were executed by a specific process. A processing trace node **520** may contain but is not limited to a correlation data field **521** identifying the portion of the monitored transaction, a topology identifier **522** identifying the process executing the portion of the monitored transaction and a payload data field **523** describing details of the execution, like performance data of executed methods, resources used for the execution, execution context data like values of method parameters or method return values.

(155) A communication trace node **530** as shown in FIG. *5e* may be used to describe a monitored communication between two processes performed during the execution of a monitored transaction. A communication trace node **530** may contain but is not limited to a communication type field **531** describing the type of the monitored communication (e.g. by the used application level protocol like HTTP, RMI or JDBC/ODBC), a sender data field **540** and a receiver data field **550**.

(156) Sender data **540** and receiver data **550** may contain but are not limited to a correlation data field identifying the portion of a monitored transaction acting as the sender of the communication activity **541** (sender data **540**) and identifying the portion of the monitored transaction acting as the receiver **551** of the communication activity (receiver data **550**), a field containing network communication endpoint data (i.e. sender and receiver network endpoint and port data) as seen from the sender side **542** and the receiver side **552** and a field containing a captured communication side as seen from the sender side **543** and the receiver side **553**. A communication trace node **530** may also contain an intermediate topology node identifier list **560** which may be populated with topology identifiers of topology entities representing identified intermediate nodes during the transaction/topology inference processing as described in FIGS. **8** to **14**.

(157) It is noteworthy that the network infrastructure may route and dispatch network traffic between a sender and a receiver process. On a logical, application level, sender and receiver are connected directly. On a network connection level, the sender process communicates with the first routing/dispatching entity and the receiver process communicates with the last routing/dispatching entity in a chain of routing/dispatching entities that connect sender and receiver. As performance problems may also arise from such routing/dispatching entities, it is important for a performance monitoring tool to detect and identify such routing/dispatching entities that are involved in the execution of a monitored transaction. Captured network level communication endpoint data as seen

on sender and receiver side may be used to detect the existence of such intermediate routing/dispatching entities.

(158) FIG. 5c shows a portion of a topology model describing a host computer node running a process that communicates with a process running on another host computer. Topology node **500a** describes a host computing system with topology identifier **501a** set to “Host A”, a not set parent topology identifier indicating that no enclosing topology entity exists for the host computing system, an entity type **503a** set to “Host” indicating that the topology node describes a host computing system and a network data field **504a** set to “Address B” indicating that the network address of the host computing system is “B”. The type-specific additional data **505a** is omitted. Topology node **500b** describes a process with topology identifier **501b** “Process C” and a parent topology identifier **502b** “Host A” indicating **560** that process “Process C” runs on “Host A”. The entity type **503b** indicates a process, and the network address data **504b** indicates that “Process C” uses port 1 to communicate with another process. Topology communication node **510a** further describes this communication and links **561** and **562** client process **500b** and server process **500c**. The client identifier **511a** identifies “Process C” as client process that initiated the communication, e.g. by sending a request. The client identifier field **511a** may, in addition, contain data describing the port (i.e. port 1) used for the communication. The server identifier field **512a** identifies “Process D” as server-side of the communication and it may also contain data describing the port that was used by the server process to receive the communication message (i.e. port 2). The processes may use TCP to transport the message, and the communication type **513** may be set accordingly. Topology node **500c** describes the receiving process by its topology identifier **501c**, its parent topology identifier **502c** which describes that “Process D” runs on host “Host E”, its type **503c** and its network address data **504c**.

(159) Typically, communication between processes relies on a stable server-side port that receives requests and uses arbitrary client-side ports. Therefore, the client-side port information is less relevant than the port on server-side. Some variant embodiments may therefore only store server-side port data.

(160) An exemplary portion of an end-to-end transaction trace is shown in FIG. 5f, which shows a communication between two monitored processes and an outgoing communication from a monitored process to a not monitored entity.

(161) Processing trace node **520a** represents transaction execution activities identified by correlation data **521a** “A”, performed on a process with topology identifier “Process E” **522a**. The monitored transaction execution on process E sends an RMI (remote method invocation) request to a process with topology identifier “Process F”. The monitored transaction execution performed on process F is represented by processing trace node **520b** with correlation data **521b** “B” and topology identifier “Process F”. Communication Trace node **530a** describes this RMI call and links **563** and **564** the corresponding transaction execution portions performed on process E and F. The communication type **531a** is set to indicate an RMI request, the correlation data **541a** of sender data **540a** identifies the transaction execution portion described by processing trace node **520a** as sender. The communication endpoint data **542a** indicates “H” as sender address (i.e. address host executing process E), and that port 3 was used on the sender side to send the message. Further, the communication endpoint data indicates “I” as receiver address and the usage of port 4 to receive the message. The communication message **543a** identifies “X” as the name of the method that is executed remotely. The correlation data **551a** of the receiver data **550a** identifies the portion of the transaction executed on process F as receiver of the message and communication endpoint data **552a** and communication message is identical to the corresponding sender data (**542a** and **543a**), which indicates that no intermediate node was involved in the transfer of the communication message.

(162) Transaction processing on process F causes the sending of a database query, which is described by communication trace node **530b**. The communication type **531b** indicates a database

query, and the sender data identifies and links **565** the transaction processing on process **F 520b** with the request. In addition, the sender data indicates that the request was sent from address “I” and port 5 to address “J” and port 6. The communication message **542b** may contain the sent database query. There may be no transaction agent **150** deployed to the process receiving the request, therefore the receiver data **550b** is empty **566**.

(163) Referring now to FIG. **6** which describes the processing of topology events as described in FIGS. **4a** to **4f** by the topology data processor **104**. FIG. **6a** shows the processing of incoming topology records and FIG. **6b** describes the identification and reporting of not monitored entities of a monitored environment using network level communication data received with topology records.

(164) Processing of topology records starts with step **600** when a new topology record is received by the topology data processor. Subsequent step **601** queries the topology model **105** if a topology node **500** with a topology identifier **501** matching the topology identifier of the received topology record exists in the topology model. In case no such topology node exists, a new one is created in step **603**, and its topology identifier **501** and payload data **505** are set to the corresponding data from the received topology record. Following step **604** sets the type-specific network data **502** (i.e. network address for host computing systems, host address and ports for processes). Step **606** sets the parent topology identifier **502** of the created topology node to the parent topology identifier reported by the received topology record. In case no parent topology identifier is reported, the topology identifier **502** is set to a value indicating no parent topology entity. Afterwards, step **607** is executed, which analyzes the communication activity data received with the topology record (i.e. network activity and endpoint data **413** received with a process topology record **410**) and creates corresponding topology communication nodes **510**.

(165) As an example, a process topology record **410** may be received that contains network activity data **413** describing communication of the process as client role with a server-side endpoint specified by a network address and a port. The network address may be used to identify the topology entity representing the host computing system on which the server process is running and the port may be used to identify the topology entity representing the server process. Port numbers are unique on an operating system and data describing the usage of ports by processes running on operating systems which may be captured by an OS agent may be used to identify the process corresponding to the port number. After the topology entities representing both client-side and server-side process are identified, a corresponding topology communication node **510** may be created. In case no topology entity for one of the communication endpoints is found, the received network endpoint data for the not resolved topology entity may be stored in the corresponding field of the created topology communication node (i.e. client topology identifier **511** or server topology identifier **512**) together with an indicator that the topology entity is not yet resolved for later resolving or other processing.

(166) The process then ends with step **608**. In case step **601** identified an existing topology node with a topology identifier **501** matching the topology identifier of the received topology record, step **609** is executed, which updates payload data and type-specific network data of the topology node found in step **601** with corresponding data from the received topology record. The process then continues with step **607**.

(167) FIG. **6b** shows an exemplary method to identify and report not monitored entities in a monitored environment. Such not monitored entities may e.g. arise when the operator of the monitored environment forgets to place agents on those entities. The process begins with step **610** when a topology communication node **510** is created and no corresponding topology node for one of the network addresses of the described communication was found. Following step **611** queries the topology model **105** for a topology communication node **510** with an adjacent unresolved endpoint matching the new received unresolved endpoint (e.g. new topology communication node has an unresolved server-side endpoint, **611** queries for existing topology communication nodes with unresolved client-side endpoint with same address as unresolved servers-side endpoint of new

topology communication node). In case no matching topology communication node is found, the process ends with step **616**. Otherwise, step **613** is executed which performs checks to determine if the unknown intermediate node (not represented in topology model but seen as a server and as client part of monitored low-level communications) is monitorable. As an example, step **613** may compare the network address of the unknown intermediate node with an address range describing the available network addresses of the monitored environment. In case the network address of the unknown intermediate node is within the address range, step **613** may indicate that the unknown node is monitorable. Otherwise, it may indicate that it is not monitorable. In case step **613** indicates that the unknown node is not monitorable, the process ends with step **616**. Otherwise, step **615** is executed which indicated the unknown intermediate node as monitoring candidate, e.g. by sending a message containing its network address to the user of the monitoring system. The process then ends with step **616**.

(168) Coming now to FIG. 7 which describes the processing of transaction trace records as described in FIGS. 4g to 4i by the transaction trace processor and the identification of finished end-to-end transaction trace records.

(169) The processing of transaction trace records, as shown in FIG. 7a starts with step **700** when a new transaction trace record is received by the transaction trace processor **102**. Following step **701** determines if the received transaction trace record is a local trace record **460**. In this case, step **702** is executed which creates a processing trace node **520** and sets its correlation data **521**, topology identifier **522** and payload data to corresponding data from the received local trace record **460**. The process then ends with step **704**. In case step **701** determines that the received transaction trace record is no local trace record, step **705** is executed which determines if the received transaction trace record is an outgoing trace record **470**. In this case, step **706** is executed which creates a new communication trace node **530** and sets its communication type according to the communication type reported by the received outgoing trace record. Following step **707** sets communication endpoint data **542**, communication message **543** and correlation data **541** of the sender data section **540** to the corresponding values of the received outgoing trace record. The process then ends with step **708**.

(170) In case step **705** determined that the received transaction trace record is no outgoing trace record, the process continues with step **709** which determines if the received transaction trace record is an incoming trace record **480**. In this case, step **710** is executed, which fetches a communication trace node **530** with correlation data **541** of the sender data **540** matching the sender correlation data **482** of the received incoming trace record **480**. Following step **711** sets correlation data **551**, communication endpoint data **552** and communication message **553** of the receiver data section **550** of the fetched communication trace node to the corresponding values of the received incoming trace record **480**. The process then ends with step **712**.

(171) The process identifying finished end-to-end transaction traces and the marking of such end-to-end transactions as ready for topology inference processing is shown in FIG. 7b. The process is executed cyclically and starts with step **720** when a specific time (5 seconds, 1 minute) since its last execution is elapsed. Following step **721** identifies those end-to-end transaction traces for which no more additional trace data can be expected. Step **721** may e.g. determine the time when the last trace record was received for a transaction. In case this time is longer than a specific time period (e.g. 1 minute, 5 minutes) in the past, the end-to-end transaction trace may be identified as finished. Following step **722** may mark finished end-to-end transaction traces identified by step **721** as ready for topology inference processing. Such marked transactions may be forwarded to the transaction/topology inference processor **109** for further analysis. The process then ends with step **723**.

(172) Coming now to FIG. 8 which conceptually describes the overall transaction/topology inference process. The process is performed for each communication trace node **530** contained in an end-to-end transaction trace record and starts with step **800** when a new finished end-to-end

transaction trace record is received by the transaction/topology inference processor **109**.

(173) Following step **801** checks if the currently processed communication trace node **530** contains sender data **540**. Sender data **540** may not be available for a call from a not monitored process that initiated the execution of a monitored transaction. In case no sender data is available, the process continues with step **806**.

(174) Otherwise, step **802** is executed, which analyzes the communication endpoint data **542** of the sender data **540** to find the topology entity representing the receiver of the communication activity as seen from the sender process. This analysis process is described in detail in FIG. **9**. Possible outcomes of this analysis are an indication that no corresponding topology entity is available in the topology model, that a network node level topology entity was found (i.e. load balancer entity or proxy appliance entity for which no monitoring data of executed processes is available) or that a process level topology entity was found (e.g. proxy process running on a monitored host).

(175) In case step **803** indicates that no corresponding topology entity was found for the receiver network endpoint reported by the sending process, step **804** is executed which updates the currently processed communication trace node and the inferred topology to represent the discovered not monitored entity. Details of the processing performed by step **804** can be found in FIG. **10**.

(176) Step **805** is executed when step **803** indicates that a corresponding topology entity was found for the receiver network endpoint reported by the sender side of the currently analyzed transaction communication. Step **805** marks the topology entity that was identified as receiver of the transaction communication as provider of the service corresponding to the type of the sent communication message (e.g. in case the type of the communication message is RMI, the identified receiving topology entity may be marked as provider of an RMI service) and enriches the communication trace node with data from the identified receiver entity (e.g. by appending the topology identifier of the identified topology entity to the intermediate topology node identifier list). A detailed description of the processing performed by step **805** can be found in FIG. **11**.

(177) Step **806** is executed after step **804** or **805**, which determines if receiver data **550** is available in the currently processed communication trace node **530**. Receiver data may not be available for transaction communication sent from instrumented processes to not instrumented receiver processes. In case of not available receiver data, the process ends with step **808**.

(178) Otherwise, the receiver side communication message **553** may be analyzed by step **807** for data indicating intermediate nodes between sender and receiver (e.g. XFF header data for HTTP based communication), corresponding topology entities for indicated intermediate nodes may be determined and the currently processed communication trace node and the inferred topology model may be updated to represent the identified intermediate nodes. A detailed description of the processing performed by step **807** can be found in FIG. **12**. The process then ends with step **808**.

(179) Coming now to FIG. **9** which describes the analysis of communication endpoint data **542** representing the network endpoints of a monitored communication as seen from the client-side. Communication endpoint data **542** consists of a network address and a port of the client/sender side and a network address and a port for the server/receiver side. The process begins with step **900** when communication endpoint data **542** is received for analysis. Following step **901** fetches the topology entity corresponding to the sender endpoint. Step **901** may first fetch the topology identifier **522** of the processing trace node **520** describing the process side processing that sent the message and then use this topology identifier to fetch the corresponding topology node describing this process.

(180) Following step **902** determines whether the sending process was executed in a container environment. Step **902** may use topology data to determine whether the sending process was executed in a container environment. OS agents **151** monitor containers and processes running on operating systems and also detect whether processes are executed directly on the operating system or in a container environment. Process id (PID) mapping data that maps operating system level PIDs to container level PIDs may be used for this determination. The OS agent may send topology

data describing which processes run in which container to the monitoring server which creates a corresponding topology model. Topology nodes **500** contain a topology identifier that uniquely identifies a topology entity. Transaction trace data contains the topology identifier of processes involved into the processing of the transaction. Therefore, the transaction trace data also contains the topology identifier of the sending process. The topology identifier of the sending process may be used to identify the topology node representing the sending process. The parent topology identifier **502** may be used to identify the topology node representing the parent topology entity of the process. In case the sending process was executed in a container, the parent topology node represents a container, the process was executed in a container, and container specific network access restrictions may apply. Data describing those container specific network access restrictions may be stored in the topology node representing the container. In case the parent topology node represents an operating system, the sending process was executed in the context of the operating system.

(181) Subsequent step **903** evaluates sender side network access and visibility restrictions arising from the container status (i.e. in case the process is executed in a container environment, other containers and processes in those other containers may be visible according to the network configuration of the container environment which are not visible from outside the container, and the container network configuration may restrict the accessibility of network components outside the container environment) or from other access and visibility restricting network configurations like virtual private clouds to identify topology entities that are reachable by the sender process.

(182) Subsequent step **904** queries those topology entities from the set of topology entities reachable by the sender process that was determined by step **903** that have a network address that matches the reported receiver address. Step **904** may further consider known traffic dispatcher components between reported sender and receiver endpoint by determining for the reported receiver network endpoint if the corresponding topology entity is a traffic dispatcher. In this case, the topology identifier of the traffic dispatcher may be added to the intermediate topology identifier list **560** of the currently processed communication trace node and the target topology nodes of the traffic dispatcher (i.e. those topology elements to which the traffic dispatcher routes incoming communication according to captured traffic mapping data) may be added to the list of reachable topology entities.

(183) Following step **905** may, in case the query performed by step **904** resulted in multiple matching topology entities use heuristics, like a match between the service requested by the sent message and services provided by the identified topology entity candidates, or matching timing between observed TCP level communication and the reported sending time of the request to determine the best matching receiver entity.

(184) Afterwards, decision step **906** checks whether a matching receiver topology entity was found. In case no matching receiver topology entity was found, step **913** is executed, which indicates that no receiver topology entity is available. The process then ends with step **914**.

(185) In case a matching receiver topology entity was found, step **907** determines if process level data is available for the identified receiver topology entity. Process level data may not be available for topology entities of virtual endpoint services like database or file storage services because such entities are typically not monitorable on a process level. If no process level data is available, step **909** is executed which indicates the found network node entity as topology entity corresponding to the receiver endpoint data.

(186) If otherwise, process level data is available for the identified receiver topology entity, step **908** is executed which uses the reported port of the receiver endpoint data to determine whether the port directs to a process or a container. Step **908** may use captured container port mapping data to determine if the reported receiver port matches a mapped container port. If no matching mapped container port is found, step **912** is executed which identifies the process corresponding to the receiver endpoint and indicates the identified process as topology entity corresponding to the



receiver endpoint data. the process then ends with step **914**.

(187) In case of a match with a mapped container port, step **911** is executed which uses the captured port mapping data to determine the corresponding port inside the container. Following step **912** then uses this corresponding container inside port to identify the receiving process inside the container and indicates this process running inside the container as topology entity corresponding to the receiver endpoint data. The process then ends with step **914**.

(188) Referring now to FIG. **10** which describes the process of handling reported network addresses for which no corresponding topology entity was found in the topology model. The process starts with step **1000** when no matching topology model entity was found for a network address reported in sender data **540** communication endpoint data **542** or receiver data **550** communication endpoint data **552**. Following step **1001** creates a new topology node **500** with an entity type **503** indicating an operating system/host computing system environment, sets the type-specific network data **504** to the reported network address for which no corresponding topology entity was found, and marks the created topology entity as unmonitored entity (e.g. by setting the type-specific additional data **505** to a value indicating an unmonitored entity).

(189) Following step **1002** adds the created topology node **500** to the topology model **105** (some variant embodiments may add topology nodes **500** describing not monitored entities to the inferred topology model **106**) and adds known horizontal and vertical relationships of the unmonitored entity to the topology model.

(190) In case a topology entity describing a virtualized host system (e.g. entity type **503** indicating a virtual host system), as e.g. reported by a cloud/virtualization agent **152**, with type-specific network data **504** matching the reported network address is available in the topology, the topology node describing the virtualized host system may be combined with the created topology node describing the unmonitored entity, e.g. by setting the parent topology identifier **502** of the created topology node to the topology identifier **501** of the topology entity describing the corresponding virtualized host system. This indicates the situation that a new virtualized host computer system was started in a monitored environment and the startup of the new virtualized host computer system was monitored by a cloud virtualization agent **152**, but there is yet no OS agent **151** installed on the virtual host computing system. Such an indication may later (e.g. step **1006**) be used to determine if a notification about the detection of a not monitored but monitorable entity of the monitored environment should be sent to the user. Step **1002** may further add a topology communication node **510** to the topology model that describes the reported transaction communication with the new unmonitored topology entity (e.g. for a monitored incoming transaction communication, from the new, unmonitored topology node to the topology node representing the known receiver of the transaction communication and for a monitored outgoing transaction communication, from the topology node representing the known sender of the transaction communication to the new, unmonitored topology node).

(191) Following decision step **1003** determines whether the currently processed communication trace node describes a transaction communication that was monitored on both sender and receiver side (i.e. a transaction agent **150** deployed to process initiating the communication and to process serving the communication), and executes step **1005** in this case, which marks the communication trace node as routed via an unknown intermediate node, e.g. by appending the topology identifier **501** of the topology node created by step **1001** to the intermediate topology node identifier list **560** of the currently processed communication trace node.

(192) In case either sender side or receiver side is not monitored by an agent, step **1004** is executed. In case the sender side is not monitored, the sender data **540** of the currently processed communication trace node may be linked to the topology node representing the new unmonitored entity, e.g. by appending a field “unmonitored topology identifier” to the sender data **540** and setting it to the topology identifier **501** of the topology node created in step **1001**. In case the receiver side is not monitored, such a field may be appended to the receiver data **550** and then also

set to the topology identifier **501** of the newly created topology node.

(193) Step **1006** is executed after step **1004** or **1005**, which evaluates criteria for the notification of the identified unmonitored entity. Typically, a notification of an unmonitored entity should only occur when evidence for the usage of the unmonitored entity in a monitored transaction is available, and the available data describing the unmonitored entity indicates that it is possible to monitor the currently unmonitored entity. Such criteria evaluated by step **1006** may contain but are not limited to a network address range, describing a set of network addresses that are within the monitored environment and determining whether the address of the unmonitored entity is within this address range, and the existence of a topology node describing a virtual host computing system with an address that matches the address of the unmonitored entity. For identified unmonitored entities on process level, those criteria may also include evaluation of a message type sent to or received from the unmonitored entity and determining, based on the message type, the type of process sending or receiving the request (e.g. process providing database services), and determining based on the type of the process whether a transaction trace agent **150** for this type of process is available.

(194) In case the notification criteria evaluated by step **1006** are fulfilled, step **1008** is executed which notifies the user of the monitoring system about the detection of an unmonitored entity in the monitored environment that is involved in the execution of a monitored transaction and that may be monitored by instrumenting it with an agent. The process then ends with step **1009**. In case the notification criteria are not fulfilled, the ends with step **1009** without sending a notification.

(195) Coming now to FIG. **11** describes the processing performed by the inference processor during in case sender data **540** is available for a currently processed communication trace node **530**, and a topology node was found in the topology model that matches the receiver address received in the communication endpoint data **542** of the sender data **540**. The process starts with step **1100** when a topology entity for a received client-side reported transaction communication receiver is was found in the topology model **105**.

(196) Following step **1101** analyzes the client-side captured data describing the communication activity, like communication type and communication message to determine the type of requested service (e.g. communication type JDBC/ODBC indicate a requested database service, communication type RMI indicate a requested remote method invocation service).

(197) Subsequent step **1102** may determine if the receiver entity is a traffic dispatcher, like a load balancer **133**. Step **1102** may e.g. use the entity type field **503** of the found topology node representing the sender side reported receiver of the communication activity for this determination.

(198) In case the receiver is no traffic dispatcher, step **1110** is executed which uses data stored in the topology node corresponding to the receiver side to enrich the receiver part of the communication trace node. As an example, the type of the communication trace node may indicate the request of a database service and the topology data describing the process that received and processed the request may identify the type, vendor and version of the database server executed by the receiver process. Step **1110** may add this data to the communication trace node to refine the transaction trace data with type, vendor, and version of the database server to which the database service request was directed.

(199) Step **1110** may, in addition, link the communication trace node describing the request with the topology node representing the process serving the request. This linkage may later be used by an automated causality and root cause detection mechanism to identify potential context conditions on the process serving the request, like a high CPU load during the time at which the request was serviced, to identify a longer than expected time required to process the request.

(200) Following step **1111** may use data extracted from the communication trace node to improve the topology model. As an example, a communication trace node **530** specifies a communication type **531**, which corresponds to a type of requested service, and the communication endpoint data specifies which address and which port was used to request this service. Topology data may only

contain data describing at which communication ports the receiving process is waiting for incoming requests, but it may not contain data describing an assignment of ports to service types. The data from the communication trace node determines the port at which the service corresponding to the communication type **531** was requested. The topology node representing the receiving process may be updated accordingly by adding descriptive data indicating that the port used by the communication trace node provides the service corresponding to the communication type.

(201) As an example, the topology model may provide data about process “A” indicating that it runs a database server (this may be determined by analyzing the name of the executable used by the process and the command line of the process) and that process “A” listens for incoming requests at ports 1,2,3 and 4. A communication trace node may specify port 2 as the port used by a request for a database service. Step **1111** may then add data to the topology model describing process “A” that indicates that port 2 provides a database service.

(202) After step **1111** the process ends with step **1112**.

(203) In case step **1102** determined that the receiver is a traffic dispatcher, step **1103** is executed, which uses reported dispatcher mapping data (see e.g. network mapping data **442** of virtual traffic dispatcher record **440**) to identify a set of potential receiver endpoints to which the dispatcher could have forwarded the request according to its mapping data. As traffic dispatchers may be chained, and the entities to which a traffic dispatcher forwards received requests may also be traffic dispatchers, step **1102** may be executed recursively in case the topology entities to which a traffic dispatcher routes requests are also traffic dispatchers, until the entities to which traffic is dispatched are no longer traffic dispatchers.

(204) Following step **1104** fetches the topology entities representing potential receiver endpoints to which dispatchers that received the transaction message may have forwarded the transaction message. Traffic dispatchers are typically configured to receive requests on one specific network endpoint and to forward received requests to one of a set of potential receivers. Traffic dispatchers may use various strategies to select a receiver for an incoming request, like a round robin strategy that sends each request to another receiver or load based strategies that monitor load parameters of potential receivers of incoming requests and send incoming requests to the receiver with the least load.

(205) Therefore, a request received by a traffic dispatcher may be forwarded to and received by one of a set of potential receivers. The traffic dispatcher mapping data captured by agents deployed to the monitored environment defines the set of potential receivers of a request, but the mapping data alone does not provide data to identify the actual receiver of the request. Therefore, step **1104** first determines those potential receivers by their network endpoint (i.e. network address and port) and determines for those network endpoints if they represent a process or a container. In case the network endpoint represents a container, captured container port mapping data may be used to identify the receiver process inside the container. In case no corresponding topology entities are found for one or more of the identified target network endpoints of the traffic dispatcher, those network endpoints with no corresponding topology entity are processed as described in FIG. **10**.

(206) Following decision step **1105** determines if data is available to identify the topology entity that received the transaction communication message described by the currently processes communication trace node. Step **1105** may determine if receiver data **550** is available in the communication trace node, which indicates that the corresponding communication message was received by a process that was instrumented with an agent and that a processing trace node **520** exists that describes the processing of the communication message on the receiving process. In case no receiver data is available, step **1105** indicates that the receiver entity could not be identified and the process continues with step **1107**. Step **1107** is executed when no data to identify the receiver process of the transaction message is available and the message was routed via one or more traffic dispatchers and any process to which the traffic dispatchers may have routed the message is a

potential receiver of the message. The currently processed communication trace node may, in this case, be adapted by appending a list containing the topology identifiers of those potential receivers which may be interpreted as “one out of n” selection by a later processing stage, like an automated causality and root cause detection mechanism. Further, the communication trace node may be adapted by appending a tree data structure describing the potential routing variants of the message via the involved traffic dispatchers. As an example, the message may first be sent to traffic dispatcher “A”, which may have forwarded it to dispatcher “B” or “C”, and dispatcher “B” may have forwarded the message to dispatcher “D” or dispatcher “E”. The root node of the tree structure may, in this case, identify dispatcher “A”, which has two child nodes “B” and “C”, and “B” would again have child nodes “D” and “E”. This tree structure describes all possible routes of the message through the involved traffic dispatchers. A later processing stage like a like an automated causality and root cause detection mechanism may use this data structure to calculate impact probabilities of an outage of one or more of those traffic dispatchers on the sender of the dispatched message.

(207) If step **1105** otherwise determines that receiver data **550** is available in the communication trace node, it may use the correlation data **551** stored in the receiver data **551** to fetch the processing trace node **520** with a matching correlation data **521** and use the topology identifier **522** of the processing trace node **520**, which identifies the topology element representing the process that received the transaction message, to indicate the topology entity representing the process as receiver of the transaction message. In this case, the process continues with step **1106**.

(208) Step **1106** may use the topology node representing the receiver of the message to determine a route of the message through the involved traffic dispatchers. In case a message is routed via multiple traffic dispatchers those traffic dispatchers typically form either a chain (one dispatcher is connected to the next dispatcher until the final receiver is reached) or in a tree (the first dispatcher dispatches to n following dispatchers and so on). In case of a chain structure, the route of the message is unambiguous. In case of a tree structure, step **1106** may use the network endpoint data of the receiver of the message to identify the dispatcher from which the message was sent to the receiver, and the mapping data of the identified dispatcher to identify the dispatcher from which it received the message and so on until the original sender of the message is reached to determine the route of the message through the dispatchers. Step **1106** may then append the topology identifiers of the identified traffic dispatchers that forwarded the message to the intermediate topology node identifier list **560** of the currently processed communication trace node **530**.

(209) Step **1108** is executed after step **1106** or **1107** and marks topology entities determined by step **1103** as potential receivers of the message as providers of the service requested by the message. As the traffic dispatchers may have delivered the message to any one of the topology entities identified by step **1103**, all those entities must be capable to provide the service requested by this message. In a more concrete example, topology data as e.g. provided by OS agents and cloud/virtualization agents may show that a traffic dispatcher “A” dispatches incoming traffic on port 2 to port 1 on host “B” and “C”, and that those ports relate to process “X” on host “B” and process “Y” on host “C”. In case a transaction message requesting a specific service, like an “HTTP” service was sent to port 2 of dispatcher “A”, the request for an “HTTP” service may have been dispatched to process “X” or “Y” on port 1. Therefore, it can be concluded that both process “X” and “Y” provide an HTTP service on port 1. After step **1108** the process ends with step **1109**.

(210) Coming now to FIG. **12** which describes the processing of the communication type **531** and the receiver side communication message **553** of the currently processed communication trace node **530** to find evidence for intermediate nodes, like proxy components, that routed the communication message from the sender to the receiver. The process starts with step **1200** when a communication trace node is received that contains receiver data **550** and that has a communication type that indicates the possibility of portions of the communication message that describe intermediate nodes. An example of such a communication type is HTTP, and examples for a portion of the

communication message describing intermediate nodes are XFF (X-Forwarded-For) or Forwarded HTTP Extension headers.

(211) Following step **1201** tries to extract data describing intermediate nodes from the captured receiver-side communication message. For a communication message following the HTTP protocol, this may be performed by extracting the values corresponding to a XFF header or headers defined by the Forwarded HTTP Extension from the communication message. Subsequent decision step **1202** checks if such data is available, e.g. by checking if the extraction performed by step **1201** provided data.

(212) In case no data describing intermediate nodes is available, step **1210** is executed, which analyzes sender and receiver side network endpoint data in conjunction with known topology data determine if intermediate nodes between sender and receiver may exist and, in case such intermediate nodes may exist, tries to identify the topology entities representing those intermediate nodes. A detailed description of the processing performed by step **1210** is found in FIG. **14**. Afterwards, the process ends with step **1211**.

(213) If step **1202** otherwise determines that data describing intermediate nodes is available, step **1203** is executed which fetches data describing the first sender of the message from the data describing intermediate nodes. An XFF header field contains a comma-separated list of the network addresses of senders from which intermediate proxy components in a proxy chain directly received a request. As an example, a sender with address “sender” may send a message to a proxy with the address “proxy1”, which forwards it to “proxy2”. The proxy component “proxy2” forwards the request to the receiver. The XFF field may be set by “proxy1” to “sender”, as it received the message from a network address “sender”. Afterwards, “proxy2” may append the network address of the sender from which it receives the message, which is “proxy1” to the XFF field, which has then the value “sender, proxy1” before the message is forwarded to the receiver. Proxies using the HTTP Forwarded Extension may add further data like the used protocol and sender/receiver ports to the request message. HTTP Forwarded Extension may also contain data describing the proxy component receiving and processing a request message. In the above example, the data generated by the HTTP Forwarded Extension would also contain data describing “proxy2”, the last proxy component in the proxy chain. It is noteworthy that providing XFF or HTTP Forwarded data is optional and there may proxy components exist in a proxy chain that do not add this data to a request. Other data, like transaction reported communication endpoint data and reported topology data, may be used to deduce the existence of such proxy components and to further identify entities of the topology model describing those proxy components.

(214) As the first entry of XFF or HTTP Forwarded data identifies the network address of the original sender of the message, and the sender data **540** of the currently processed communication trace node **530** contains data describing the network address of the sender, step **1204** may compare both network addresses to determine if a not reported intermediate node may have processed the transaction message before the first proxy component that added proxy identification data (i.e. XFF or HTTP Forwarded data) to the request.

(215) A sender address extracted from the proxy identification data matching the sender address extracted from the sender data **540** indicates that there exists no not reported intermediate node between sender and first reporting proxy (i.e. first proxy that adds XFF or HTTP Forwarded data to the request). In case the extracted sender addresses do not match, step **1204** may use the sender side captured receiver data of the message (i.e. the portion of the communication endpoint data **542** captured by the sender side agent), the sender address extracted from the captured receiver-side communication message and data from the topology model to identify not reported intermediate nodes between the sender and the first reported proxy entity. First, topology nodes corresponding to the sender side captured receiver network address and to the sender address extracted from the communication message may be fetched and appended to the intermediate topology node identifier list **560** of the currently processed communication trace node, as those topology nodes represent the

network entities to which the sender sent the message and from which the first reported proxy element received the message.

(216) Afterwards, step **1204** may analyze the topology recorded communication paths between the topology node with the network address to which the sender sent the message and the topology node with the network address from which the first reported proxy received the message. In case there is one unambiguous communication path between both topology nodes, the topology nodes involved in this communication path may also be appended to the intermediate topology node identifier **560**. As an example, the sender side reported receiver may have network address “A”, the network address of the sender reported by the first reported proxy may be “X”. Topology data shows that “A” communicates with node “B” and node “C”, node “C” communicates only with “D”, node “D” communicates with node “X”, and there is no other communication path between “A” and “X” than via “C” and “D”. In this case, nodes “C” and “D” may be inserted into the intermediate topology node identifier list **560** between nodes “A” and “X”.

(217) Following step **1205** checks if the data describing intermediate nodes extracted by step **1201** contains data describing a next intermediate sender node. In case no data for a next intermediate sender node is available, step **1212** is executed which checks if data for the last intermediate node is available. Such data may be available for proxies reporting according to the HTTP Forwarded Extension standard, as this defines, next to a “Forwarded For” header which contains data describing the network node from which a proxy received a message, a “Forwarded By” header which contains data describing the proxy that received the message. Step **1212** may e.g. check if a “Forwarded By” header is available in the server-side captured communication message and may in this case execute step **1213**, which analyzes the network address provided by the “Forwarded By” header in combination with the receiver side reported sender address (i.e. sender network address in communication endpoint data **552** of the receiver data **550**) and the topology model to infer the existence of unreported intermediate node and to identify topology nodes corresponding to those unreported intermediate node as described in step **1204**.

(218) After step **1212** or in case no data describing the last intermediate node is available, the process ends with step **1214**.

(219) In case step **1205** finds data for a next intermediate sender node, step **1206** is executed which fetches the topology node corresponding to the network address of the next intermediate node. In case no matching topology node is found, step **1208** is executed which handles the unavailable topology entity as described in FIG. **10**. If otherwise a matching topology was found by step **1206**, step **1209** is executed which uses data describing the communication activity and the topology node, like reported port mapping data, to find a more specific topology node. As an example, the matching topology node may identify a host computer system and step **1209** may identify the topology node describing the process running on the host computer that provides the proxy service that was used to forward the message. Afterwards, step **1209** may update the currently processed communication trace node and the inferred topology model **106** according to its findings. A detailed description of the processing performed by step **1209** can be found in FIG. **13**. After step **1208** or **1209** the process continues with step **1206** which determines if data describing another next sender node is available.

(220) Coming now to FIG. **13** which describes the process that uses topology and transaction-related monitoring data to determine for a network address of an intermediate node (e.g. proxy or load balancer node) if a process running on the host computing system identified by the network address can be identified that provides the used intermediate node service.

(221) The process starts with step **1300** when a topology node representing an intermediate node reported by a transaction trace was identified. Following decision step **1301** determines whether the currently analyzed topology node represents a dedicated traffic dispatcher entity. A dedicated traffic dispatcher entity is a network node that does not operate a conventional operating system on which processes may be executed, but a virtualized or concrete entity restricted to a traffic dispatching

functionality. Examples are Amazon Elastic Load Balancer® for virtualized entities or the F5 BIG-IP® load balancer for dedicated hardware appliances. In case the currently analyzed topology node represents a dedicated traffic dispatcher entity, no process level data is available and the process ends with step **1302**. Step **1301** may use the entity type **503** of the topology node to determine if it represents a dedicated traffic dispatcher entity.

(222) In case the currently analyzed topology node is no dedicated traffic dispatcher entity but a host computing system, step **1303** is executed which determines if in addition to the network address data of the intermediate node, also port data is available that describes the port which was used by the intermediate node to send or receive a message. Such port data may be available if intermediate nodes report forwarding activity according to the specification of the HTTP Forwarded Extension.

(223) In case port data is available, step **1304** is executed which uses the port data to identify a process level entity executed on the previously identified host computing system. Step **1304** may also consider whether the provided port identifies a process or a container and in case a container is identified, use captured container port mapping data to identify the corresponding process executed inside the container. After the topology node representing the process that provides the traffic dispatching or proxy service was identified by step **1304**, step **1304** may further analyze the type-specific additional data **505** describing the process to verify if the process provides a traffic dispatcher or proxy service. Step **1304** may analyze captured command line, executable name or configuration data of the process. An example for software-based traffic dispatcher or proxy services that are executed as processes on an operating system is the HAProxy software product.

(224) In case no port data is available, step **1305** is executed which uses topology reported process and container data and reported process level network mapping data as e.g. extracted from configuration data of identified traffic dispatcher and proxy processes in combination with heuristics to identify the process level entity that provided the used traffic dispatching/proxy service. Step **1305** may determine the number of traffic dispatcher or proxy processes running on the host computing system. In case there is only one such process, this process may be notified as used traffic dispatching or proxy process. In case there are more than one matching processes, monitored network communication activities of those processes may be analyzed to select those processes that had a network connection with an adjacent node in the traffic dispatcher/proxy chain during the execution of the request. In case there is only one such traffic dispatcher/proxy process, it may be notified by step **1305** as used traffic dispatching or proxy process. Otherwise, step **1305** may indicate that no process level entity was found.

(225) Step **1306** is executed after step **1304** or **1305** and checks whether a traffic dispatching or proxy process was identified. In case none was identified, the process ends with step **1309**. Otherwise, step **1307** is executed which updates the currently processed communication trace node **560** to describe the usage of the identified traffic dispatcher/proxy process. Step **1307** may e.g. replace in the intermediate topology identifier list **560** the topology identifier identifying the host computing system executing the traffic dispatcher/proxy process with the topology identifier of the identified traffic dispatcher/proxy process. Afterwards, step **1308** is executed which may, in case a process was identified for which it was unknown that it provides a traffic dispatcher/proxy service, update the inferred topology model **106** to describe this new finding. The process then ends with step **1309**.

(226) Referring now to FIG. **14**, which describes the processing of sender and receiver endpoint data as e.g. available in communication endpoint data **542** of sender data **540** and in communication endpoint data **552** of receiver data **550** of a communication trace node **530**. The process starts with step **1400** when a pair of sender and receiver network endpoint data is received. Subsequent step **1401** checks if the sender and receiver endpoint data are equal, which indicates that the receiver endpoint as seen from the sender side is equal to the receiver endpoint reported by the receiver and that the sender endpoint as seen from the receiver is equal to the sender endpoint

reported by the receiver. In case sender and receiver endpoint data are equal, the process ends with step **1403**.

(227) Otherwise, step **1402** is executed which determines if sender and receiver endpoint data indicate a single intermediate node. A single intermediate node is indicated when the receiver endpoint reported by the sender side is equal to the sender endpoint reported by the receiver side. This indicates that the sender sent its message to a specific network node that is different from the receiver and the receiver received the message from this specific network node. Therefore, this situation indicates that only one intermediate node forwarded the message from the sender to the receiver.

(228) Step **1405** is executed in case step **1402** indicates a single intermediate node. Step **1405** queries the reported topology model **105** for a network node entity (e.g. host computing system running an operating system or dedicated traffic dispatcher) with a network address matching the network address of the single intermediate node identified by step **1402**. Step **1405** may also consider sender side network restrictions, like virtual private cloud or container network configurations affecting the network nodes that are visible for the sender and only consider those network nodes that are visible for the query of the topology model.

(229) In case no topology node with matching network address is found, step **1407** is executed which handles a reported network address for which no topology node was found as described in FIG. **10**.

(230) In case a matching topology node was found, step **1408** is executed which updates the currently processing communication trace node **530** by adding the topology identifier of the identified topology node to the intermediate topology node identifier list **560**. In addition, processing as described in FIG. **13** may be executed to identify the intermediate node on a process level. In addition, the inferred topology model may be updated to describe that the identified intermediate topology node provides a traffic dispatching/proxy service. After step **1407** or **1408** the process ends with step **1409**.

(231) In case step **1404** determined that multiple intermediate nodes routed the message from the sender to the receiver (i.e. receiver network address seen from the sender side is not equal to sender network address seen from the receiver side), step **1410** is executed which uses the reported topology model to determine an unambiguous path between the reported endpoints, by e.g. analyzing communication activity described in the reported topology model to determine if there exists only one communication path between sender and receiver. In case such an unambiguous path exists, step **1412** is executed which updates the currently analyzed transaction communication node **530**, e.g. by appending the topology identifiers of the topology nodes on the identified path to its intermediate topology node identifier list **560**. In addition, step **1412** may update the inferred topology model **106** by adding data indicating that all topology entities on the identified path provide a traffic dispatching/proxy service. The process ends with step **1409** after step **1412** or if no unambiguous path was found by step **1410**.

(232) Coming now to FIG. **15**, which conceptually describes a possible usage transaction trace and topology data that was improved by a transaction/topology inference processing to improve the quality of an automated causality and root cause detection process for identified performance, functionality and resource usage anomalies.

(233) The process starts with step **1500** when continuous monitoring of transaction execution and architecture topology of a computing environment, like e.g. a data center, is started. A monitoring server **101** may be installed and started during step **1501**. Afterwards, agents, like transaction agents **150**, OS agents **151** and cloud/virtualization agents **152** are installed to components of the computing environment that should be monitored. Some portions of step **1501** may be performed manually, other portions, like the injection of transaction agents **150** into processes may be performed automatically. The installed agents may be configured to send monitoring data to the monitoring server. Afterwards, the agents may be started.



(234) The monitoring server **101** may in parallel receive transaction trace data **1502** and topology data **1504**. Received topology data may be used to update **1505** the reported topology model **105**, while received transaction trace data may be used in step **1503** by a transaction/topology inference processor **109** to infer unknown topology elements, to improve the inferred topology model **106** and to enhance the received transaction data as described before.

(235) Following step **1507** analyzes performance and functional behavior data of incoming transaction trace data and resource usage and functionality of topology entities to identify anomalies and locate those anomalies in the topology model.

(236) Subsequent step **1508** uses inference enhanced transaction trace data and data from the inferred topology model to estimate causal relationships between identified anomalies and to estimate a set of anomalies that may be the root cause for other anomalies. The data gained by the transaction/topology inference process helps to improve this causal correlation of anomalies because it extends the knowledge base used by the analysis. As an example, anomaly detection may detect that the execution time of a certain type of a transaction increased, but there is no evidence explaining this change on monitored processes involved into the execution of the transaction. Transaction/topology inference may have revealed that messages used to execute the transaction are routed via a certain proxy process, and resource usage monitoring data of the host computing system running the proxy process shows that CPU usage of this host computing system has peaked during the execution of the transaction that showed the performance anomaly. Automated root cause detection may use this additional data to identify the CPU peak at the host computing system as the root cause for the performance degradation of the transaction executions.

(237) The techniques described herein may be implemented by one or more computer programs executed by one or more processors. The computer programs include processor-executable instructions that are stored on a non-transitory tangible computer readable medium. The computer programs may also include stored data. Non-limiting examples of the non-transitory tangible computer readable medium are nonvolatile memory, magnetic storage, and optical storage.

(238) Some portions of the above description present the techniques described herein in terms of algorithms and symbolic representations of operations on information. These algorithmic descriptions and representations are the means used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. These operations, while described functionally or logically, are understood to be implemented by computer programs. Furthermore, it has also proven convenient at times to refer to these arrangements of operations as modules or by functional names, without loss of generality.

(239) Unless specifically stated otherwise as apparent from the above discussion, it is appreciated that throughout the description, discussions utilizing terms such as “processing” or “computing” or “calculating” or “determining” or “displaying” or the like, refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system memories or registers or other such information storage, transmission or display devices.

(240) Certain aspects of the described techniques include process steps and instructions described herein in the form of an algorithm. It should be noted that the described process steps and instructions could be embodied in software, firmware or hardware, and when embodied in software, could be downloaded to reside on and be operated from different platforms used by real time network operating systems.

(241) The present disclosure also relates to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, or it may comprise a computer selectively activated or reconfigured by a computer program stored on a computer readable medium that can be accessed by the computer. Such a computer program may be stored in a tangible computer readable storage medium, such as, but is not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, magnetic-optical disks, read-only memories

(ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, application specific integrated circuits (ASICs), or any type of media suitable for storing electronic instructions, and each coupled to a computer system bus. Furthermore, the computers referred to in the specification may include a single processor or may be architectures employing multiple processor designs for increased computing capability.

(242) The algorithms and operations presented herein are not inherently related to any particular computer or other apparatus. Various systems may also be used with programs in accordance with the teachings herein, or it may prove convenient to construct more specialized apparatuses to perform the required method steps. The required structure for a variety of these systems will be apparent to those of skill in the art, along with equivalent variations. In addition, the present disclosure is not described with reference to any particular programming language. It is appreciated that a variety of programming languages may be used to implement the teachings of the present disclosure as described herein.

(243) The foregoing description of the embodiments has been provided for purposes of illustration and description. It is not intended to be exhaustive or to limit the disclosure. Individual elements or features of a particular embodiment are generally not limited to that particular embodiment, but, where applicable, are interchangeable and can be used in a selected embodiment, even if not specifically shown or described. The same may also be varied in many ways. Such variations are not to be regarded as a departure from the disclosure, and all such modifications are intended to be included within the scope of the disclosure.

## Claims

1. A computer-implemented method for monitoring transactions executing in a distributed computing environment, comprising: monitoring, by a given agent, a communication message exchanged between two processes executing in the distributed computing environment, wherein the given agent is deployed in one of the two processes; capturing, by the given agent, monitoring data for the communication message, where in the monitoring data includes at least one of network address data for process sending the communication message and network address data for process receiving the communication message; sending, by the given agent, the monitoring data to a monitoring server, wherein the monitoring server is remotely located from a computing device hosting the given agent; extracting, by the monitoring server, the network address data for the process sending or receiving the communication message from the monitoring data; determining, by the monitoring server, whether the extracted network address data for the process sending or receiving the communication message is stored at the monitoring server; and extracting, by the monitoring server, data describing a network component from the monitoring data and storing the data together with the extracted network address data for the process sending or receiving the communication message on the monitoring server, wherein the extracting and storing of the data is in response to a determination that the extracted network address data for the process sending or receiving the communication message is not stored at the monitoring server; wherein the given agent is deployed in the process sending the communication message and the monitoring data further includes an indicator of communication protocol used to exchange the communication message between the two processes and further comprises: extracting, by the monitoring server, the indicator of the communication protocol from the monitoring data; determining, by the monitoring server, a type of process for the process receiving the communication message from the indicator of the communication protocol; extracting, by the monitoring server, the network address data for the process receiving the communication message; and storing, by the monitoring server, the type of process for the process receiving the communication message together with the extracted network address data for the process receiving the communication message on the monitoring server, wherein the storing is in response to a determination that the extracted network address data for the

process receiving the communication message is not stored at the monitoring server.

2. The computer-implemented method of claim 1 wherein the indicator of the communication protocol indicates a database query, such that the type of process for the process receiving the communication message is a database.

3. The computer-implemented method of claim 1 further comprising receiving topology data from an operating system (OS) agent instrumented in an operating system of a computing device, wherein the topology data identifies the operating system of the computing device, identifies the computing device, identifies processes handled by the operating system of the computing device, and includes network mapping data.

4. The computer-implemented method of claim 3 wherein the network mapping data is further defined as container port mapping data, and wherein the container port mapping data specifies a relationship between a port and a process executing in a container.

5. The computer-implemented method of claim 1 further comprising: capturing, by the given agent, traffic mapping configuration data from a traffic dispatching component; sending, by the given agent, the traffic mapping configuration data to the monitoring server; and translating, by the monitoring server, network address data for process at which the given agent is deployed on into network address data seen by another process using the traffic mapping configuration data.

6. A computer-implemented method for monitoring transactions executing in a distributed computing environment, comprising: monitoring, by a given agent, a communication message exchanged between two processes executing in the distributed computing environment, wherein the given agent is deployed in one of the two processes; capturing, by the given agent, monitoring data for the communication message, where in the monitoring data includes at least one of network address data for process sending the communication message and network address data for process receiving the communication message; sending, by the given agent, the monitoring data to a monitoring server, wherein the monitoring server is remotely located from a computing device hosting the given agent; extracting, by the monitoring server, the network address data for the process sending or receiving the communication message from the monitoring data; determining, by the monitoring server, whether the extracted network address data for the process sending or receiving the communication message is stored at the monitoring server; and extracting, by the monitoring server, data describing a network component from the monitoring data and storing the data together with the extracted network address data for the process sending or receiving the communication message on the monitoring server, wherein the extracting and storing of the data is in response to a determination that the extracted network address data for the process sending or receiving the communication message is not stored at the monitoring server; wherein the given agent is deployed in the process receiving the communication message and the monitoring data further includes additional data that is added to the communication message by an intermediate component and further comprises; extracting, by the monitoring server, the additional data from the monitoring data, wherein the extracted additional data includes a type for the intermediate component and network address data for the intermediate component; and storing, by the monitoring server, the extracted additional data together with the extracted network address data for the process receiving the communication message on the monitoring server, wherein the storing is in response to a determination that the extracted network address data for the process receiving the communication message is not stored at the monitoring server.

7. The computer-implemented method of claim 6 wherein the type for the intermediate component indicates a proxy server.

8. A computer-implemented method for monitoring transactions executing in a distributed computing environment, comprising: monitoring, by a given agent, a communication message exchanged between two processes executing in the distributed computing environment, wherein the given agent is deployed in one of the two processes; capturing, by the given agent, monitoring data for the communication message, where in the monitoring data includes at least one of network

address data for process sending the communication message and network address data for process receiving the communication message; sending, by the given agent, the monitoring data to a monitoring server, wherein the monitoring server is remotely located from a computing device hosting the given agent; extracting, by the monitoring server, the network address data for the process sending or receiving the communication message from the monitoring data; determining, by the monitoring server, whether the extracted network address data for the process sending or receiving the communication message is stored at the monitoring server; and extracting, by the monitoring server, data describing a network component from the monitoring data and storing the data together with the extracted network address data for the process sending or receiving the communication message on the monitoring server, wherein the extracting and storing of the data is in response to a determination that the extracted network address data for the process sending or receiving the communication message is not stored at the monitoring server; wherein one agent is deployed in the process sending the communication message and another agent is deployed in the process receiving the communication message and further comprises; capturing, by the one agent, network address data for process sending the communication message and network address data for process receiving the communication message; sending, by the one agent, the network address data for process sending the communication message and the network address data for process receiving the communication message to the monitoring server; capturing, by the another agent, the network address data for process sending the communication message and the network address data for process receiving the communication message; sending, by the another agent, the network address data for process sending the communication message and the network address data for process receiving the communication message to the monitoring server; comparing, by the monitoring server, the network address data for process sending the communication message captured by the one agent to the network address data for process sending the communication message captured by the another agent, and comparing, by the monitoring server, the network address data for process receiving the communication message captured by the one agent to the network address data for process receiving the communication message captured by the another agent; and indicating, by the monitoring server, delivery of the communication message from the process sending the communication message directly to the process receiving the communication message, wherein the indication is in response to a match between the network address data for process sending the communication message as captured by the one agent and the network address data for process sending the communication message as captured by the another agent, and the network address data for process receiving the communication message as captured by the one agent and the network address data for process receiving the communication message as captured by the another agent.

9. The computer-implemented method of claim 8 further comprising indicating, by the monitoring server, that the communication message was delivered via exactly one intermediate node, wherein the indication is in response to a match between the network address data for process receiving the communication message as captured by the one agent and the network address data for process receiving the communication message as captured by the another agent.

10. The computer-implemented method of claim 8 further comprising indicating, by the monitoring server, the communication message was delivered by multiple intermediate nodes in response to a mismatch between the network address data for process sending the communication message as captured by the one agent and the network address data for process sending the communication message as captured by the another agent, and the network address data for process receiving the communication message as captured by the one agent and the network address data for process receiving the communication message as captured by the another agent.

11. A system for monitoring transactions executing in a distributed computing environment, comprising: a given agent deployed in one of two processes executing in the distributed computing environment, wherein the given agent is configured to: monitor a communication message

exchanged between the two processes, capture monitoring data for the communication message, and send the monitoring data to a monitoring server, and wherein the monitoring data includes at least one of network address data for process sending the communication message and network address data for process receiving the communication message; and a non-transitory tangible computer readable medium storing computer-executable instructions that, upon execution of the computer-executable instructions by a processor of the monitoring server, cause the monitoring server to: extract the network address data for the process sending or receiving the communication message from the monitoring data; determine whether the extracted network address data for the process sending or receiving the communication message is stored at the monitoring server; and extract data describing a network component from the monitoring data and store the data together with the extracted network address data for the process sending or receiving the communication message on the monitoring server, wherein the extracting and storing of the data is in response to a determination that the extracted network address data for the process sending or receiving the communication message is not stored at the monitoring server; wherein the given agent is deployed in the process sending the communication message and the monitoring data further includes an indicator of communication protocol used to exchange the communication message between the two processes and the computer-executable instructions further cause the monitoring server to: extract the indicator of the communication protocol from the monitoring data; determine a type of process for the process receiving the communication message from the indicator of the communication protocol; extract the network address data for the process receiving the communication message; and store the type of process for the process receiving the communication message together with the extracted network address data for the process receiving the communication message on the monitoring server, wherein the storing is in response to a determination that the extracted network address data for the process receiving the communication message is not stored at the monitoring server.

12. The system of claim 11 wherein the indicator of the communication protocol indicates a database query, such that the type of process for the process receiving the communication message is a database.

13. A system for monitoring transactions executing in a distributed computing environment, comprising: a given agent deployed in one of two processes executing in the distributed computing environment, wherein the given agent is configured to: monitor a communication message exchanged between the two processes, capture monitoring data for the communication message, and send the monitoring data to a monitoring server, and wherein the monitoring data includes at least one of network address data for process sending the communication message and network address data for process receiving the communication message; and a non-transitory tangible computer readable medium storing computer-executable instructions that, upon execution of the computer-executable instructions by a processor of the monitoring server, cause the monitoring server to: extract the network address data for the process sending or receiving the communication message from the monitoring data; determine whether the extracted network address data for the process sending or receiving the communication message is stored at the monitoring server; and extract data describing a network component from the monitoring data and store the data together with the extracted network address data for the process sending or receiving the communication message on the monitoring server, wherein the extracting and storing of the data is in response to a determination that the extracted network address data for the process sending or receiving the communication message is not stored at the monitoring server; wherein the given agent is deployed in the process receiving the communication message and the monitoring data further includes additional data that is added to the communication message by an intermediate component and the computer-executable instructions further cause the monitoring server to: extract the additional data from the monitoring data, where in the extracted additional data includes a type for the intermediate component and network address data for the intermediate component; and store the extracted

additional data together with the extracted network address data for the process receiving the communication message on the monitoring server, wherein the storing is in response to a determination that the extracted network address data for the process receiving the communication message is not stored at the monitoring server.

14. The system of claim 13 wherein the type for the intermediate component indicates a proxy server.

15. A system for monitoring transactions executing in a distributed computing environment, comprising: a given agent deployed in one of two processes executing in the distributed computing environment, wherein the given agent is configured to: monitor a communication message exchanged between the two processes, capture monitoring data for the communication message, and send the monitoring data to a monitoring server, and wherein the monitoring data includes at least one of network address data for process sending the communication message and network address data for process receiving the communication message; and a non-transitory tangible computer readable medium storing computer-executable instructions that, upon execution of the computer-executable instructions by a processor of the monitoring server, cause the monitoring server to: extract the network address data for the process sending or receiving the communication message from the monitoring data; determine whether the extracted network address data for the process sending or receiving the communication message is stored at the monitoring server; and extract data describing a network component from the monitoring data and store the data together with the extracted network address data for the process sending or receiving the communication message on the monitoring server, wherein the extracting and storing of the data is in response to a determination that the extracted network address data for the process sending or receiving the communication message is not stored at the monitoring server; wherein one agent is deployed in the process sending the communication message and another agent is deployed in the process receiving the communication message, wherein the one agent captures network address data for process sending the communication message and network address data for process receiving the communication message and sends the network address data for process sending the communication message and the network address data for process receiving the communication message to the monitoring server; and wherein the another agent captures the network address data for process sending the communication message and the network address data for process receiving the communication message and sends the network address data for process sending the communication message and the network address data for process receiving the communication message to the monitoring server.

16. The system of claim 15 wherein the computer-executable instructions further cause the monitoring server to: compare the network address data for process sending the communication message captured by the one agent to the network address data for process sending the communication message captured by the another agent; compare the network address data for process receiving the communication message captured by the one agent to the network address data for process receiving the communication message captured by the another agent; and indicate delivery of the communication message from the process sending the communication message directly to the process receiving the communication message, where in the indication is in response to a match between the network address data for process sending the communication message as captured by the one agent and the network address data for process sending the communication message as captured by the another agent, and the network address data for process receiving the communication message as captured by the one agent and the network address data for process receiving the communication message as captured by the another agent.

17. The system of claim 15 wherein the computer-executable instructions further cause the monitoring server to indicate that the communication message was delivered via exactly one intermediate node, and wherein the indication is in response to a match between the network address data for process receiving the communication message as captured by the one agent and the

network address data for process receiving the communication message as captured by the another agent.

18. The system of claim 15 wherein the computer-executable instructions further cause the monitoring server to indicate the communication message was delivered by multiple intermediate nodes in response to a mismatch between the network address data for process sending the communication message as captured by the one agent and the network address data for process sending the communication message as captured by the another agent, and the network address data for process receiving the communication message as captured by the one agent and the network address data for process receiving the communication message as captured by the another agent.

---