

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386759
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Roweth; Duncan et al.

Algorithms for use of load information from neighboring nodes in adaptive routing

Abstract

Systems and methods are provided for passing data amongst a plurality of switches having a plurality of links attached between the plurality of switches. At a switch, a plurality of load signals are received from a plurality of neighboring switches. Each of the plurality of load signals are made up of a set of values indicative of a load at each of the plurality of neighboring switches providing the load signal. Each value within the set of values provides an indication for each link of the plurality of links attached thereto as to whether the link is busy or quiet. Based upon the plurality of load signals, an output link for routing a received packet is selected, and the received packet is routed via the selected output link.

Inventors: Roweth; Duncan (Bristol, GB), Froese; Edwin L. (Burnaby, CA), Kopnick; Joseph G. (Chippewa Falls, WI), Kopser; Andrew S. (Seattle, WA), Alverson; Robert (Seattle, WA)

Applicant: Hewlett Packard Enterprise Development LP (Spring, TX)

Family ID: 73458112

Assignee: Hewlett Packard Enterprise Development LP (Spring, TX)

Appl. No.: 18/479755

Filed: October 02, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240039836 A1	Feb. 01, 2024

Related U.S. Application Data

Publication Classification

Int. Cl.: **H04L45/125** (20220101); **G06F9/50** (20060101); **G06F9/54** (20060101); **G06F12/0862** (20160101); **G06F12/1036** (20160101); **G06F12/1045** (20160101); **G06F13/14** (20060101); **G06F13/16** (20060101); **G06F13/28** (20060101); **G06F13/38** (20060101); **G06F13/40** (20060101); **G06F13/42** (20060101); **G06F15/173** (20060101); **H04L1/00** (20060101); **H04L43/0876** (20220101); **H04L43/10** (20220101); **H04L45/00** (20220101); **H04L45/02** (20220101); **H04L45/021** (20220101); **H04L45/028** (20220101); **H04L45/12** (20220101); **H04L45/122** (20220101); **H04L45/16** (20220101); **H04L45/24** (20220101); **H04L45/42** (20220101); **H04L45/745** (20220101); **H04L47/10** (20220101); **H04L47/11** (20220101); **H04L47/12** (20220101); **H04L47/122** (20220101); **H04L47/20** (20220101); **H04L47/22** (20220101); **H04L47/24** (20220101); **H04L47/2441** (20220101); **H04L47/2466** (20220101); **H04L47/2483** (20220101); **H04L47/30** (20220101); **H04L47/32** (20220101); **H04L47/34** (20220101); **H04L47/52** (20220101); **H04L47/62** (20220101); **H04L47/625** (20220101); **H04L47/6275** (20220101); **H04L47/629** (20220101); **H04L47/76** (20220101); **H04L47/762** (20220101); **H04L47/78** (20220101); **H04L47/80** (20220101); **H04L49/00** (20220101); **H04L49/101** (20220101); **H04L49/15** (20220101); **H04L49/90** (20220101); **H04L49/9005** (20220101); **H04L49/9047** (20220101); **H04L67/1097** (20220101); **H04L69/22** (20220101); **H04L69/40** (20220101); **H04L45/28** (20220101); **H04L45/7453** (20220101); **H04L69/28** (20220101)

U.S. Cl.:

CPC **G06F13/1642** (20130101); **G06F9/505** (20130101); **G06F9/546** (20130101); **G06F12/0862** (20130101); **G06F12/1036** (20130101); **G06F12/1063** (20130101); **G06F13/14** (20130101); **G06F13/16** (20130101); **G06F13/1673** (20130101); **G06F13/28** (20130101); **G06F13/385** (20130101); **G06F13/4022** (20130101); **G06F13/4068** (20130101); **G06F13/4221** (20130101); **G06F15/17331** (20130101); **H04L1/0083** (20130101); **H04L43/0876** (20130101); **H04L43/10** (20130101); **H04L45/02** (20130101); **H04L45/021** (20130101); **H04L45/028** (20130101); **H04L45/122** (20130101); **H04L45/123** (20130101); **H04L45/125** (20130101); **H04L45/16** (20130101); **H04L45/20** (20130101); **H04L45/22** (20130101); **H04L45/24** (20130101); **H04L45/38** (20130101); **H04L45/42** (20130101); **H04L45/46** (20130101); **H04L45/566** (20130101); **H04L45/70** (20130101); **H04L45/745** (20130101); **H04L47/11** (20130101); **H04L47/12** (20130101); **H04L47/122** (20130101); **H04L47/18** (20130101); **H04L47/20** (20130101); **H04L47/22** (20130101); **H04L47/24** (20130101); **H04L47/2441** (20130101); **H04L47/2466** (20130101); **H04L47/2483** (20130101); **H04L47/30** (20130101); **H04L47/32** (20130101); **H04L47/323** (20130101); **H04L47/34** (20130101); **H04L47/39** (20130101); **H04L47/52** (20130101); **H04L47/621** (20130101); **H04L47/6235** (20130101); **H04L47/626** (20130101); **H04L47/6275** (20130101); **H04L47/629** (20130101); **H04L47/76** (20130101); **H04L47/762** (20130101); **H04L47/781** (20130101); **H04L47/80** (20130101); **H04L49/101** (20130101);

H04L49/15 (20130101); **H04L49/30** (20130101); **H04L49/3009** (20130101);
H04L49/3018 (20130101); **H04L49/3027** (20130101); **H04L49/90** (20130101);
H04L49/9005 (20130101); **H04L49/9021** (20130101); **H04L49/9036** (20130101);
H04L49/9047 (20130101); **H04L67/1097** (20130101); **H04L69/22** (20130101);
H04L69/40 (20130101); G06F13/1689 (20130101); G06F2212/50 (20130101);
G06F2213/0026 (20130101); G06F2213/3808 (20130101); H04L45/28 (20130101);
H04L45/7453 (20130101); H04L69/28 (20130101)

Field of Classification Search

CPC: H04L (45/28); H04L (43/0876); H04L (43/00); H04L (43/10); H04L (45/123); H04L (45/38); H04L (45/70); H04L (47/00); H04L (47/11); H04L (47/12); H04L (47/122); H04L (47/20); H04L (47/22); H04L (47/30); H04L (49/101); H04L (49/15); H04L (49/50); G06F (9/505)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2012/0314577	12/2011	Sun	370/235	H04L 47/26
2019/0018743	12/2018	Takaoaka	N/A	G06F 11/0727

Primary Examiner: Kamara; Mohamed A

Attorney, Agent or Firm: Dickinson Wright

Background/Summary

RELATED APPLICATIONS (1) This application is a continuation of and claims the priority of U.S. patent application Ser. No. 17/594,736, filed Oct. 28, 2021, which is a 371 National Stage Entry of PCT/US2020/024129, filed on Mar. 23, 2020, which claims the benefit of U.S. Provisional Patent Application No. 62/852,273, filed May 23, 2019, U.S. Provisional Patent Application No. 62/852,203, filed May 23, 2019, and U.S. Provisional Patent Application No. 62/852,289, filed May 23, 2019, the disclosures of which are incorporated herein by reference in their entirety for all purposes.

DESCRIPTION OF RELATED ART

(1) As network-enabled devices and applications become progressively more ubiquitous, various types of traffic as well as the ever-increasing network load continue to demand more performance from the underlying network architecture. For example, applications such as high-performance computing (HPC), media streaming, and Internet of Things (IOT) can generate different types of traffic with distinctive characteristics. As a result, in addition to conventional network performance metrics such as bandwidth and delay, network architects continue to face challenges such as scalability, versatility, and efficiency.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present disclosure, in accordance with one or more various embodiments, is described in detail with reference to the following figures. The figures are provided for purposes of illustration only and merely depict typical or example embodiments.
- (2) FIG. 1 illustrates an example network in which various embodiments may be implemented.
- (3) FIG. 2 illustrates an example switching system that facilitates flow channels in accordance with various embodiments.
- (4) FIG. 3A illustrates crossbars implemented within an example crossbar switch in accordance with various embodiments.
- (5) FIG. 3B illustrates an example tile matrix corresponding to ports of the example edge switching system of FIG. 2 in accordance with various embodiments.
- (6) FIG. 3C illustrates an example tile making up the tile matrix of FIG. 3B in accordance with various embodiments.
- (7) FIG. 3D illustrates an example age queue implementation in accordance with various embodiments.
- (8) FIG. 4A and FIG. 4B are block diagrams of an example FRF component implemented at each port of the example edge switching system of FIG. 2.
- (9) FIG. 5 illustrates an example of route selection in accordance with various embodiments.
- (10) FIG. 6 illustrates an example of local and global load value updating and use in accordance with various embodiments.
- (11) FIG. 7A illustrates an example of average load-based routing.
- (12) FIG. 7B illustrates an example of neighboring switch load-based adaptive routing in accordance with various embodiments.
- (13) FIG. 8 illustrates an example computing component capable of executing instructions for load-based adaptive routing in accordance with one embodiment.
- (14) FIG. 9 is an example computing component in which various embodiments described herein may be implemented.
- (15) The figures are not exhaustive and do not limit the present disclosure to the precise form disclosed.

DETAILED DESCRIPTION

- (16) The present disclosure describes systems and methods that can accommodate exascale computing, e.g., perform data-intensive tasks such as simulations, data analytics, artificial intelligence workloads at exascale speeds. In particular, an HPC network or interconnect fabric is provided that may be Ethernet-compatible, able to connect to third-party data storage, and can be built using a switch component that is extremely high bandwidth, e.g., on the order to 12.8 Tb/s/dir per switch with, e.g., 64 200 Gbps ports that support large network creation with very low diameter (e.g., only three network hops). Moreover, low latency can be achieved through novel congestion control mechanisms, adaptive routing, and the use of traffic classes that allows for flexibility in terms of bandwidth shaping, priority, and routing policy.
- (17) When passing information from switch to switch, it is necessary to make instantaneous decisions about the next hop on the path. Decisions are taken using information derived from local state and information that is communicated from neighboring switches. Use of information from many different sources allows for more accurate/effective decisions. This includes information from neighbors.
- (18) Prior systems carried information on average load from switch to switch. That said, more detailed information from related or neighboring switches is more helpful. Accordingly, in some embodiments, a set of values can include information indicative of the status of output ports of related/neighboring switches. By passing this set of values, much better routing decisions can be made. In one example, a flag is passed back from neighboring switches, with the flag having one

bit for each output port. For example, with a switch having 64 outputs, a 64 bit flag would be transmitted. This is much more accurate than simply passing a global average for neighboring ports.

(19) FIG. 1 shows an example network **100** comprising a plurality of switches, which can also be referred to as a “switch fabric.” As illustrated in FIG. 1, network **100** can include switches **102**, **104**, **106**, **108**, and **110**. Each switch can have a unique address or ID within switch fabric **100**. Various types of devices and networks can be coupled to a switch fabric. For example, a storage array **112** can be coupled to switch fabric **100** via switch **110**; an InfiniBand (IB) based HPC network **114** can be coupled to switch fabric **100** via switch **108**; a number of end hosts, such as host **116**, can be coupled to switch fabric **100** via switch **104**; and an IP/Ethernet network **118** can be coupled to switch fabric **100** via switch **102**. For example, a switch, such as switch **102** may receive 802.3 frames (including the encapsulated IP payload) by way of Ethernet devices, such as network interface cards (NICs), switches, routers, or gateways. IPv4 or IPv6 packets, frames formatted specifically for network **100**, etc. may also be received, transported through the switch fabric **100**, to another switch, e.g., switch **110**. Thus, network **100** is capable of handling multiple types of traffic simultaneously. In general, a switch can have edge ports and fabric ports. An edge port can couple to a device that is external to the fabric. A fabric port can couple to another switch within the fabric via a fabric link.

(20) Typically, traffic can be injected into switch fabric **100** via an ingress port of an edge switch, and leave switch fabric **100** via an egress port of another (or the same) edge switch. An ingress edge switch can group injected data packets into flows, which can be identified by flow ID's. The concept of a flow is not limited to a particular protocol or layer (such as layer-2 or layer-3 in the Open System Interface (OSI) reference model). For example, a flow can be mapped to traffic with a particular source Ethernet address, traffic between a source IP address and destination IP address, traffic corresponding to a TCP or UDP port/IP 5-tuple (source and destination IP addresses, source and destination TCP or UDP port numbers, and IP protocol number), or traffic produced by a process or thread running on an end host. In other words, a flow can be configured to map to data between any physical or logic entities. The configuration of this mapping can be done remotely or locally at the ingress edge switch.

(21) Upon receiving injected data packets, the ingress edge switch can assign a flow ID to the flow. This flow ID can be included in a special header, which the ingress edge switch can use to encapsulate the injected packets. Furthermore, the ingress edge switch can also inspect the original header fields of an injected packet to determine the appropriate egress edge switch's address, and include this address as a destination address in the encapsulation header. Note that the flow ID can be a locally significant value specific to a link, and this value can be unique only to a particular input port on a switch. When the packet is forwarded to the next-hop switch, the packet enters another link, and the flow-ID can be updated accordingly. As the packets of a flow traverses multiple links and switches, the flow IDs corresponding to this flow can form a unique chain. That is, at every switch, before a packet leaves the switch, the packet's flow ID can be updated to a flow ID used by the outgoing link. This up-stream-to-down-stream one-to-one mapping between flow ID's can begin at the ingress edge switch and end at the egress edge switch. Because the flow ID's only need to be unique within an incoming link, a switch can accommodate a large number of flows. For example, if a flow ID is 11 bits long, an input port can support up to 2048 flows. Furthermore, the match pattern (one or more header fields of a packet) used to map to a flow can include a greater number of bits. For instance, a 32-bit long match pattern, which can include multiple fields in a packet header, can map up to 2^{32} different header field patterns. If a fabric has N ingress edge ports, a total number of $N \times 2^{32}$ identifiable flows can be supported.

(22) A switch can assign every flow a separate, dedicated input queue. This configuration allows the switch to monitor and manage the level of congestion of individual flows, and prevent head-of-

queue blocking which could occur if shared buffer were used for multiple flows. When a packet is delivered to the destination egress switch, the egress switch can generate and send back an acknowledgement (ACK) in the upstream direction along the same data path to the ingress edge switch. As this ACK packet traverses the same data path, the switches along the path can obtain the state information associated with the delivery of the corresponding flow by monitoring the amount of outstanding, unacknowledged data. This state information can then be used to perform flow-specific traffic management to ensure the health of the entire network and fair treatment of the flows. As explained in more detail below, this per-flow queuing, combined with flow-specific delivery acknowledgements, can allow the switch fabric to implement effective, fast, and accurate congestion control. In turn, the switch fabric can deliver traffic with significantly improved network utilization without suffering from congestion.

(23) Flows can be set up and released dynamically, or “on the fly,” based on demand. Specifically, a flow can be set up (e.g., the flow-ID to packet header mapping is established) by an ingress edge switch when a data packet arrives at the switch and no flow ID has been previously assigned to this packet. As this packet travels through the network, flow IDs can be assigned along every switch the packet traverses, and a chain of flow IDs can be established from ingress to egress. Subsequent packets belonging to the same flow can use the same flow IDs along the data path. When packets are delivered to the destination egress switch and ACK packets are received by the switches along the data path, each switch can update its state information with respect to the amount of outstanding, unacknowledged data for this flow. When a switch's input queue for this flow is empty and there is no more unacknowledged data, the switch can release the flow ID (i.e., release this flow channel) and re-use the flow-ID for other flows. This data-driven dynamic flow setup and teardown mechanism can obviate the need for centralized flow management, and allows the network to respond quickly to traffic pattern changes.

(24) Note that the network architecture described herein is different from software-defined networks (SDN's), which typically uses the OpenFlow protocol. In SDN, switches are configured by a central network controller, and packets are forwarded based one or more fields in the layer-2 (data link layer, such as Ethernet), layer-3 (network layer, such as IP), or layer-4 (transport layer, such as TCP or UDP) headers. In SDN such header-field lookup is performed at every switch in the network, and there is no fast flow ID-based forwarding as is done in the networks described herein. Furthermore, because the OpenFlow header-field lookup is done using ternary content-addressable memory (TCAM), the cost of such lookups can be high. Also, because the header-field mapping configuration is done by the central controller, the setup and tear-down of each mapping relationship is relatively slow and could require a fair amount of control traffic. As a result, an SDN network's response to various network situations, such as congestion, can be slow. In contrast, in the network described herein, the flows can be set up and torn down dynamically based on traffic demand; and packets can be forwarded by a fixed-length flow ID. In other words, flow channels can be data driven and managed (i.e., set up, monitored, and torn down) in a distributed manner, without the intervention of a central controller. Furthermore, the flow ID-based forwarding can reduce the amount of TCAM space used and as a result a much greater number of flows can be accommodated.

(25) Referring to the example shown in FIG. 1, suppose that storage array **112** is to send data using TCP/IP to host **116**. During operation, storage array **112** can send the first packet with host **116**'s IP address as the destination address and a predetermined TCP port specified in the TCP header. When this packet reaches switch **110**, the packet processor at the input port of switch **110** can identify a TCP/IP 5-tuple of this packet. The packet processor of switch **110** can also determine that this 5-tuple currently is not mapped to any flow ID, and can allocate a new flow ID to this 5-tuple. Furthermore, switch **110** can determine the egress switch, which is switch **104**, for this packet based on the destination (i.e., host **116**'s) IP address (assuming switch **110** has knowledge that host **116** is coupled to switch **104**). Subsequently, switch **110** can encapsulate the received packet with a

fabric header that indicates the newly assigned flow ID and switch **104**'s fabric address. Switch **110** can then schedule the encapsulated packet to be forwarded toward switch **104** based on a fabric forwarding table, which can be computed by all the switches in fabric **100** using a routing algorithm such as link state or distance vector.

(26) Note that the operations described above can be performed substantially at line speed with little buffering and delay when the first packet is received. After the first packet is processed and scheduled for transmission, subsequent packets from the same flow can be processed by switch **110** even faster because the same flow ID is used. In addition, the design of the flow channels can be such that the allocation, matching, and deallocation of flow channels can have substantially the same cost. For example, a conditional allocation of a flow channel based on a lookup match and a separate, independent deallocation of another flow channel can be performed concurrently in nearly every clock cycle. This means that generating and controlling the flow channels can add nearly no additional overhead to the regular forwarding of packets. The congestion control mechanism, on the other hand, can improve the performance of some applications by more than three orders of magnitude.

(27) At each switch along the data path (which includes switches **110**, **106**, and **104**), a dedicated input buffer can be provided for this flow, and the amount of transmitted but unacknowledged data can be tracked. When the first packet reaches switch **104**, switch **104** can determine that the destination fabric address in the packet's fabric header matches its own address. In response, switch **104** can decapsulate the packet from the fabric header, and forward the decapsulated packet to host **116**. Furthermore, switch **104** can generate an ACK packet and send this ACK packet back to switch **110**. As this ACK packet traverses the same data path, switches **106** and **110** can each update their own state information for the unacknowledged data for this flow.

(28) In general, congestion within a network can cause the network buffers to fill. When a network buffer is full, the traffic trying to pass through the buffer ideally should be slowed down or stopped. Otherwise, the buffer could overflow and packets could be dropped. In conventional networks, congestion control is typically done end-to-end at the edge. The core of the network is assumed to function only as “dumb pipes,” the main purpose of which is to forward traffic. Such network design often suffers from slow responses to congestions, because congestion information often cannot be sent to the edge devices quickly, and the resulting action taken by the edge devices cannot always be effective in removing the congestion. This slow response in turn limits the utilization of the network, because to keep the network free of congestion the network operator often needs to limit the total amount of traffic injected into the network. Furthermore, end-to-end congestion control usually is only effective provided that the network is not already congested. Once the network is heavily congested, end-to-end congestion control would not work, because the congestion notification messages can be congested themselves (unless a separate control-plane network that is different from the data-plane network is used for sending congestion control messages).

(29) In contrast, the flow channels can prevent such congestion from growing within the switch fabric. The flow channel mechanism can recognize when a flow is experiencing some degree of congestion, and in response can slow down or stop new packets of the same flow from entering the fabric. In turn, these new packets can be buffered in a flow channel queue on the edge port and are only allowed into the fabric when packets for the same flow leave the fabric at the destination edge port. This process can limit the total buffering requirements of this flow within the fabric to an amount that would not cause the fabric buffers to become too full.

(30) With flow channels, the switches have a reasonably accurate state information on the amount of outstanding in-transit data within the fabric. This state information can be aggregated for all the flows on an ingress edge port. This means that the total amount of data injected by an ingress edge port can be known. Consequently, the flow channel mechanism can set a limit on the total amount of data in the fabric. When all edge ports apply this limit action, the total amount of packet data in

the entire fabric can be well controlled, which in turn can prevent the entire fabric from being saturated. The flow channels can also slow the progress of an individual congested flow within the fabric without slowing down other flows. This feature can keep packets away from a congestion hot spot while preventing buffers from becoming full and ensuring free buffer space for unrelated traffic.

(31) FIG. 2 illustrates an example switch **202** (which may be an embodiment of any one or more of switches **102**, **104**, **106**, **108**, and **110**) that may be used to create a switch fabric, e.g., switch fabric **100** of FIG. 1. In this example, a switch **202** can include a number of communication ports, such as port **220**. Each port can include a transmitter and a receiver. Switch **202** can also include a processor **204**, a storage device **206**, and a flow channel switching logic block **208**. Flow channel switching module **208** can be coupled to all the communication ports and can further include a crossbar switch **210**, an EFCT logic block **212**, an IFCT logic block **214**, and an OFCT logic block **216**.

(32) Crossbar switch **210** includes crossbars which can be configured to forward data packets and control packets (such as ACK packets) among the communication ports. EFCT logic block **212** can process packets received from an edge link and map the received packets to respective flows based on one or more header fields in the packets. In addition, EFCT logic block **212** can assemble FGFC Ethernet frames, which can be communicated to an end host to control the amount of data injected by individual processes or threads. IFCT logic block **214** can include the IFCT, and perform various flow control methods in response to control packets, such as endpoint-congestion-notification ACKs and fabric-link credit-based flow control ACKs. OFCT logic block **216** can include a memory unit that stores the OFCT and communicate with another switch's IFCT logic block to update a packet's flow ID when the packet is forwarded to a next-hop switch.

(33) In one embodiment, switch **202** is an application-specific integrated circuit (ASIC) that can provide **64** network ports that can operate at either 100 Gbps or 200 Gbps for an aggregate throughput of 12.8 Tbps. Each network edge port may be able to support IEEE 802.3 Ethernet, and Optimized-IP based protocols as well as Portals, an enhanced frame format that provides support for higher rates of small messages. Ethernet frames can be bridged based on their L2 address or they can be routed based on their L3 (1Pv4/1Pv6) address. Optimized-IP frames may only have an L3 (1Pv4/1Pv6) header, and are routed. Specialized NIC support can be used for the Portals enhanced frame format, and can map directly onto the fabric format of network **100**, e.g., a fabric format that provides certain control and status fields to support a multi-chip fabric when switches/switch chips, such as switches **102**, **104**, **106**, **108**, and **110** are connected and communicate with each other. As alluded to above, a congestion control mechanism based on flow channels can be used by such switches, and can also achieve high transmission rates for small packets (e.g., more than 1.2 billion packets per second per port) to accommodate the needs of HPC applications.

(34) Switch **202** can provide system-wide Quality of Service (QoS) classes, along with the ability to control how network bandwidth is allocated to different classes of traffic, and to different classes of applications, where a single privileged application may access more than one class of traffic. Where there is contention for network bandwidth, arbiters select packets to forward based on their traffic class and the credits available to that class. Network **100** can support minimum and maximum bandwidths for each traffic class. If a class does not use its minimum bandwidth, other classes may use the unused bandwidth, but no class can get more than its maximum allocated bandwidth. The ability to manage bandwidth provides the opportunity to dedicate network resources, as well as CPUs and memory bandwidth to a particular application.

(35) In addition to support for QoS classes, switch **202** effectuates flow channel-based congestion control, and can reduce the number of network hops, e.g., in a network having a dragonfly topology, from five network hops to three. The design of switch **202**, described in greater detail below, can reduce network cost and power consumption, and may further facilitate use of

innovative adaptive routing algorithms that improve application performance. A fabric created by a plurality of switches, such as a plurality of switches **202** may also be used in constructing Fat-Tree networks, for example when building a storage subsystem for integration with third-party networks and software. Further still, the use of switch **202** enables fine-grain adaptive routing while maintaining ordered packet delivery. In some embodiments, switch **202** may be configured to send the header of a packet from an input port to an output port before the full data payload arrives, thereby allowing output port load metrics to reflect future loads, thereby improving adaptive routing decisions made by switch **202**.

(36) Crossbar switch **210** may comprise separate, distributed crossbars routing data/data elements between input and output ports. In some embodiments, and as illustrated in FIG. 3A, there are five distributed crossbars including a request crossbar **210a**, a grant crossbar **210b**, credit crossbar **210c**, an Ack crossbar **210d**, and a data crossbar **210e** between input port **220b** and output port **220c**.

(37) Request crossbar **210a** is used to send requests from an input to a targeted output age queue. Grant crossbar **210b** is used to return a grant back to the input to satisfy a request. In particular, grant crossbar **210b** returns a pointer indicating where a packet is within an input buffer. It should be noted that a grant is returned when there is space in the output for the corresponding packet. Grant crossbar **201b** may also optionally return a credit for requested space in the output. It should be noted that grants are returned when there is a landing spot for a packet at the output, e.g., an output port **220c**, so packets cannot be blocked (though they can face transient contention for resources).

(38) It should be understood that in accordance with various embodiments, a credit protocol may be used to guarantee that there is a landing space for a request at the output. Accordingly, a credit crossbar **210c** may be used to return credit for requested space in the output.

(39) A data crossbar **210d** is used to move granted packets from an input buffer to a targeted output buffer. An Ack crossbar **210e** is used to propagate Ack packets from output ports **220c** to input ports **220b**. Acks are steered in accordance with a state kept in an output flow channel table.

(40) It should be understood that data crossbar **210d** moves multi-clock packets with both headers and data, while the other four crossbars (request crossbar **210a**, grant crossbar **210b**, credit crossbar **210c**, and Ack crossbar **210e**) move only single-clock packet headers. All five crossbars use the same architecture with row buses and column buses within an 8×4 matrix of 32 dual-port tiles (as described below).

(41) Referring back to FIG. 2, switch **202** may have a plurality of transmit/receive ports, e.g., port **220**. The plurality of ports may be structured in a tile matrix. FIG. 3B illustrates an example of such a tile matrix **300**. In one embodiment, tile matrix **300** comprises 32 tiles, each comprising two ports used to implement the crossbar switching between ports, and to provide the following: a serializer/de-serializer (SERDES) interface between the core of switch **202** and external high speed serial signals for driving the signals off switch **202**; a media access control (MAC) sub-layer interface to the physical coding sublayer (PCS); a PCS interface between the SERDES and the Ethernet MAC function; a link level retry (LLR) function that operates on a per packet basis and uses ordered sets to deliver initialization sequences, Acks, and Nacks; and an Ingress Transforms block for converting between different frame fabric formats. Each tile contains a crossbar switch such as crossbar switch **210** for each of the crossbars (**210a-210e**). As will be discussed in greater detail below, routing in the switch fabric may be controlled by a fabric routing function (FRF) implemented in switch **202**, where a separate instance of the FRF component (FIGS. 4A, 4B) may be implemented within input logic for each port of switch **202**. As noted above, each port includes an instance of the FRF component, but for ease of reference/illustration clarity, only two example FRF instances (**400a**, **400b**) are labeled.

(42) Each crossbar switch **210** has sixteen inputs **220b**, one for each port in its row, and eight outputs **220c**, one for each port in its column. Row buses can be driven from each source in a row to all eight crossbars in that row (one-to-all). Arbitration can be performed at the crossbar from the

sixteen row buses in that row to the eight column buses in a given column. Buffering can be provided at each 16×8 crossbar for each of the row buses in order to absorb packets during times when there is contention for a column bus. In some embodiments, a non-jumbo packet is kept off a row bus unless there is room for the entire packet in the targeted crossbar input buffer. Due to area constraints, jumbo packets are allowed to go even if there is not sufficient space (crossbar input buffer only sized to sink a non-jumbo packet) with the row bus being blocked until the packet wins arbitration and space is freed as it is moved onto a column bus.

(43) Column buses are driven from a given crossbar to each destination port within a column (all-to-all). Each destination may have another level of arbitration between the column buses from the four rows. With sixteen row buses driving eight crossbars, each feeding eight column buses, there is a $4 \times$ speedup between rows and columns. Each row has identical connections with the one-to-all row bus connections for a single row shown in row buses. Each tile will have a one (request, grant, credit) or a two (data, ack) clock delay per tile depending on the crossbar. This gives a maximum seven or fourteen clock delay to get between the leftmost and rightmost columns. Credit returns routed through credit crossbar **210c** may have a one clock delay per tile, and therefore, can take a maximum of seven clocks to complete transmission.

(44) It should be noted that each column may have identical connections with the all-to-all column bus connections for a single column, and there may be a two clock delay per tile, resulting in a six clock delay to get from the top row to the bottom row. It should also be understood that both row and column buses both use the aforementioned credit-based protocol to determine when they are able to send. In the case of row buses, the source port maintains credit counts for the input buffers of the crossbars within that row. For the data crossbar, care is needed to determine when a packet is allowed to go on a row bus. If grants targeting a particular crossbar input buffer all go through a single queue, space for the packet at the head of the queue is required before starting the packet transfer. If the grants are distributed across multiple queues, in order to prevent small packets from locking out large packets, a packet transfer does not start unless there is space for an entire max sized packet in the buffer. In this way, once a packet transfer on a row bus starts, it will not stop until the entire packet has been transferred. Accordingly, crossbar input buffers are configured to be large enough to handle the maximum packet size plus additional space to cover the worst case round trip (packet send to credit return). This will not be the case for jumbo packets. To save on buffering area, the crossbar input buffers are only deep enough to handle a non-jumbo sized MTU (1500 bytes) with a jumbo packet being allowed to block a row bus while waiting to gain access to the targeted column bus.

(45) For column buses, each crossbar maintains credit counts for the input buffers at each destination port in that column. Unlike row buses, there is no requirement that credits be available for a maximum-sized packet before starting transfer of that packet on a column bus. Individual words of the packet will move as credits become available. Therefore, the input buffer at the destination for each column bus needs to only be big enough to cover the worst case round trip (packet to credit).

(46) FIG. 3C illustrates, in greater detail, an example implementation of two ports, e.g., ports 0 and 1, handled by tile 1, along with crossbar **220a** comprising a set of row buses and column channels with per tile crossbars. In this way, every port has its own row bus, which communicates across its row, and every tile has the aforementioned 16×8 crossbar, which is used to do corner turns, and a set of eight column channels that feed up to the eight ports that are contained in that column. In other words, each crossbar switch **210** has sixteen row bus input buffers and eight possible destinations. For example, for data to travel from, e.g., input port 17 to output port 52, data is routed along a row bus from input port 17, traverses a local crossbar which is a 16 to 8 arbitration, and then traverses up a column channel to output port 52. In terms of the total routing through all the set of distributed crossbars, there is four times more internal bandwidth than there is external bandwidth, resulting in an ability to keep up with ingress when routing nearly any arbitrary

permutation of traffic through switch **202**.

(47) A fair round-robin arbitration may be used between the sixteen sources for each destination. For the data crossbar **210d**, once a source wins arbitration, it keeps control of the destination column bus until the entire packet has been sent. Each output grants a limited amount of packet payload so it is expected that contention for a given column bus should be fairly limited when larger packets are involved. Because of this, a round-robin arbitration is expected to be sufficient even with possibly large differences in packet size among requesters.

(48) Parts of switch **202** associated with output functions generally operate on frames within the switch fabric format, and have a fabric header, even, for example, for a frame arriving and leaving on an Ethernet port within a single switch **202**.

(49) Age queue output control is responsible for accepting requests from all of the input port, e.g., input ports **220b**, via request crossbar **210a**, buffering the requests, arbitrating between them by traffic class using a traffic shaper, and passing the requests to the OFCT **216** to be granted via grant crossbar **210b**. Age queue buffering is managed to allow each input to have enough space to flow while also allowing an input with multiple flows targeting a given output to take more space. In particular, an age queue space is managed by output control. The age queue/output control may also be responsible for managing access to the link either using credit-based flow control for a connected input buffer or pause-based flow control for non-fabric links. When a packet is released by the age queue, it is committed to being put on the link. Additionally the age queue has a path allowing packets initiated on a given port e.g., one of input ports **220b** (such as maintenance or reduction packets), to arbitrate for resources on the given port.

(50) Requests come into the output control block via a column bus from each row of matrix **30**. Each column bus feeds an independent FIFO (e.g., first-in-first-out shift register or buffer) with space in the FIFO managed via credits. The FIFOs may be sized (24 deep) to cover a round-trip plus additional space to allow requests to be moved out of the crossbars **210a-210e** and prevent head-of-line blocking. Prior to writing into a FIFO, a request may be checked for a valid error correcting code (ECC). If the ECC check has either a multi bit error (MBE) or a single bit error (SBE) in the destination field (i.e. it has been routed to the wrong port), the request is considered to be an invalid request, and is discarded with an error being flagged.

(51) Least recently used (LRU) arbitration may be performed between column bus FIFOs to choose which FIFO gets forwarded to age queue management. As requests are removed from each FIFO, credits are returned to the corresponding crossbar. The row with which an incoming column bus corresponds can be dependent both on where in the matrix the tile is located, and as well as which half of the tile the block is in.

(52) The output buffer (OBUF) makes requests to the output control block for sending reduction and maintenance packets across a link. These requests may be given the highest priority. A FIFO with 8 locations can be used to buffer these reduction/maintenance packet requests while they wait for resources. Reduction packets need not use flow channels, and maintenance packets may use loopback to create a flow so that checking for flow channel availability or flowing through the OFCT to create a grant is not needed. Reduction and maintenance packets also need not use any space in the output buffer so that no check of space is required. Rather, a check for the link partner input buffer may be performed. If allowed, a shaping queue (SQ) or virtual channel (VC) can be granted, blocking any grants from the age queue path from being granted during that cycle.

(53) The size of the next request to be processed from the output buffer is checked against max_frame_size. If it exceeds this setting, the request is not processed and an error flag is set. This will result in the output buffer request path being blocked until a warm reset is performed. The error flag will stay set until the reset is done. The condition can also be released by increasing the setting of max_frame_size to a value above the size of the stuck output buffer request. The size used in the comparison may be the size indicated in the output buffer request (which may include a 4-byte frame checksum (FCS) used on the wire).

(54) Each input may be given the same fixed allocation of age queue space. This age queue space is large enough to reserve a location for each SQ/VC with enough additional space to cover a request/credit round-trip. It is up to the input to manage the space it is given across its SQs/VCS. This allocation (fixed_al/oc) is programmable via a control and status register (CSR) in each input queue (INQ), and can be, e.g., in the range of 64-96 locations. The remaining age queue space (8K-64*fixed_al/oc) may be shared space that is available to all inputs. The shared space can be managed by the output with it moving incoming requests from static to shared space as they arrive if there is room in the shared space, subject to per-input limits. When moving a request to the shared space, a credit is returned, e.g., immediately, via credit crossbar **210c**, with the request marked in the age queue as being in the shared space.

(55) When a request is granted, if it is marked as using the shared space, the shared space is credited. If it is not marked as using shared space, the request is considered to have used the static space, and a credit is returned to the input with the grant.

(56) Due to conflicts in credit crossbar **210c**, it is possible that credits may not be sent every clock period. Accordingly, a FIFO provides buffering for these transient disruptions. Space in this FIFO is required before taking a request from the request crossbar. A FIFO with a depth of 32 locations can be used to limit the chances of it ever backing up into request crossbar **210a**. The shared space may have limits for how much space any input (from an input port **220b**) can take. These limits can be set as a percentage of the available space. For instance, if the limit is set to 50%, if one input port is active, it has access to 50% of the buffer space, with two active input ports, each gets 37.5% $((\text{space_used_by_I} + \text{pace_left} * 0.5) / 2 = (50\% + 50\% * 0.5) / 2)$, with three active input ports, each gets 29.2% $((\text{space_used_by_2} + \text{space_left} * 0.5) / 3 = (75\% + 25\% * 0.5) / 3)$, and so on. Additionally, the total space used by the active input ports can be limited to the given total (50%, 75%, 87.5%). Thus, the space allocated to each of input port **220b** may vary dynamically by how many inputs ports are currently active. The addition of an active input port causes other active inputs ports to give up their space which is then taken by the new input.

(57) Given that division is not something easily done in hardware, the aforementioned age queue credit management function can be implemented as a lookup table **310** with 64 entries **312**. The number of inputs currently active in the age queues **320** indexes **315** the lookup table **310**. The values **314** in the lookup table **310** reflect the limit of the number of shared space locations any input can take along with the total space they can consume as a whole. Thus, it is up to software to program the values **314** in the lookup table **310** according to how much total shared space there is and what percentage each input port is allowed to take. As more input ports **220b** become active, each input port **220b** is allowed less space, and the total space available increases. Incoming requests from input ports **220b** that are above this limit, or in total, exceed the total space limit, are not allowed to take more shared space. In order to track the number of active input ports **220b** in the age queues, a set of 64 counters **316** (one for each input port) is used. These count up when a request is put in the age queues **320** and count down as they are taken out (i.e., granted). A count of the number of non-zero counts **319** is used as an index into the lookup table **310**. In addition, in order to manage the shared space, an additional set of 64 counters **318** may be used to track the current usage of the shared space by each input. There may also be a single counter **334** that can be used to track overall shared space usage. These counters are compared against the current quotas to determine if a request is allowed to use the shared space or not. Counters **316**, **318** can be, e.g., 13 bits wide, to provide sufficient coverage of the maximum value of an object that may be somewhat less than 8K.

(58) Age queues **320** may use a single storage RAM **321** that has 8K locations in it. These locations can be dynamically allocated to 32 separate queues (one for each SQ/VC) with each consisting of a linked-list of locations within the storage RAM **321**. This gives each SQ/VC the ability to take more space as needed.

(59) An age queue **320** can be created with a front pointer **322** pointing to the front of the queue,

and a next pointer **324** for each location pointing the next item in the queue. The last location in the queue may be indicated by a back pointer **326**. Items are taken from the front of the queue and inserted at the back of the queue. In addition to the above data structures, each queue has a FIFO **328** of entries at its head. These FIFOs **328** may ensure that a queue can sustain a request every clock with a multi-clock read access time from the request RAM **321**. When a new request arrives, if the head FIFO **328** for that queue is not full, it bypasses the request RAM **321**, and can be written directly into the head FIFO **328**. Once requests for a given age queue are being written to the request RAM **321**, subsequent requests are also written to the request RAM **321** to maintain order. The bypass path can be used again once there are no more requests for that age queue in the request RAM **321**, and there is room in the head FIFO **328**. When a request is read from a head FIFO **328**, and there are corresponding requests queued in the request RAM **321**, a dequeue is initiated. One head FIFO **328** may be read at a time, such that a single dequeue operation can be initiated each clock period. Logic may be included to handle the various race conditions between an ongoing or imminent enqueue operation and a head FIFO **328** being read.

(60) The aforementioned ECC protection used in the age queue RAM **321** can be extended to the FIFOs **328** to protect the data path flops. The resulting structure may include 8K flops (32 queues $\times$ 5 deep $\times$ SQ-bits wide). When generating the ECC, the age queue number can be included in the calculation (but not stored) as an extra check of the free list management. When the ECC is checked, the request can be considered to be in error if there is an MBE or there is an SBE in the queue number bits.

(61) A free list RAM can be a simple FIFO which is initialized with pointers to all 8K entries whenever a reset is performed. A count can be maintained to keep track of how many entries are valid within the free list. As entries are taken, they are popped off the front of the FIFO and used. As entries are returned, they are pushed onto the back of the FIFO. Some number of entries, e.g., three entries,) at the head of the free list can be kept in flops so they are available for quick access. As with the head FIFOs for the age queues, ECC is carried through the flops to provide protection. The resulting structure may have minimal flops (57=3 deep $\times$ 19-bits wide).

(62) In order to support full performance for small packets, age queues support both an enqueue and a dequeue every clock period. The operations across the data structures for an enqueue operation are discussed below, and can differ depending on whether the queue being written is empty or not.

(63) In some cases, a simultaneous enqueue and dequeue to a specific queue is easily handled as they are using and updating separate fields. Some specialized scenarios may arise, e.g., when a dequeue operation empties the age queue. In order to handle this scenario, a dequeue occurs first logically, followed by an enqueue operation. Accordingly, an empty flag is seen as being set when the queue is emptied by the dequeue operation, and then cleared due to the enqueue operation.

(64) The arbitration alluded to above can be performed among requests that are permitted to be granted subject to input buffer management, output buffer management, and flow channel quotas. Arbitration can also be halted if there are no credits for the OFCT input FIFO. In some embodiments, arbitration may be performed at two levels. First, traffic shaping arbitration can be used to arbitrate between the SQs. A Deficit Round-robin arbitration can be used to arbitrate between VCs within a given SQ. Traffic shaping arbitration may use a series of token buckets to control the bandwidth of each SQ as follows: eight leaf buckets, one for each SQ; four branch buckets; and a single head bucket.

(65) Arbitration can be divided into three groups with a first group having the highest priority, followed by a second group, which in turn is followed by a third group. For the first and second groups, arbitration may be handled in the same way among eligible SQs. A $\times 8$ round-robin arbitration can be performed between the SQs for each of the eight priority levels (eight parallel round-robin arbitrations). A fixed arbitration can be performed between priority levels. For example, group 3 arbitration has no priorities, and therefore is simply a single $\times 8$ round-robin

arbitration.

(66) For arbitration in the first group, the priority for each comes from the setting in the leaf buckets. For arbitration in the second group, priority comes from the setting in the branches of the leaf buckets. In all cases, the buckets which are checked to be eligible for that group, are also the buckets from which packet size tokens are obtained if that request wins arbitration.

(67) Regarding age queue **320** selection, packets can be classified in order to select the SQ to which their request is forwarded. This allows traffic associated with an application to be shaped differently from traffic originating from a different application or a different traffic class. This can be useful on the edge ports which connect to a NIC in that the applications will have been configured to use a share of the resources on the node, and similarly will be granted a proportion of the network bandwidth. In accordance with one embodiment, this classification is performed by classifying the packets into a traffic class identifier (FTAG), e.g., a 4-bit code that is part of the fabric frame header, and a VLAN ID (VNI) as the packet ingresses into the fabric. The FTAG and VNI may then be used as the packet egresses the fabric to select the shaping queue.

(68) A configuration register can be used to map FTAGs to SQs. This configuration matches the corresponding configuration in the in queue. When the output buffer requests or returns link partner credits, it converts a given FTAG to an SQ. For packet injection, the FTAG is found in R_TF_OBUF_CFG_PFG_TX_CTRL. For test generation, the FTAG is found in the test control register. When the reduction engine (RED) requests a credit return, the FTAG is found in ret_cdt/tag. When a reduction frame is removed from the output stream and link partner credits need to be returned, the FTAG is found in the frame header.

(69) Regarding the SQs discussed herein, each age queue **320** may have 32 SQs that are addressed by {SQ, VC}. The 3-bit SQ **330** can be considered a shaping function, and the VC selects one of four queues within that shaping function. For Ethernet egress (edge) ports, the VC is not needed for deadlock avoidance. Accordingly, all 32 SQs **330** can be available. In such a scenario, the SQ **330** can be selected by adding the SQ base from R_TF_OBUF_CFG_FTAG_SQ_MAP to the lower bits of the VNI. The 5-bit sum defines the {SQ,VC} to send to the age queue. It should be noted that when injecting frames on an egress port, a VNI is not available, and therefore, an SQ base can be directly used. For fabric links, the SQ **330** is taken from the upper three bits of the SQ base. The VC can be taken from the frame header when returning credits for reduction frames, or from the appropriate control CSR (R_TF_OBUF_CFG_TEST_CTRL or R_TF_OBUF_CFG_PFG_TX_CTRL) when injecting frames.

(70) A link partner input buffer management can depend on the type of device to which the link is attached. Devices such as switch **202** may use credit-based flow control where each credit represents a cell of storage in the input buffer. Other devices may use standard Ethernet pause or priority pause-based flow control. Requests which are marked to terminate locally (lac term set) need not consider link partner input buffer flow control and need not update any associated counters. Link partner space need not be considered when the link is in the draining state.

(71) For credit-based flow control, the link partner input buffer can be divided into eight buffer classes. Each SQ **330** can be assigned to one of these 8 buffer classes. Credits are maintained for each of the buffer classes with each credit representing 32 bytes of storage in the link partner input buffer. In order to allow credit-based flow control to work with various devices (switch, enhanced NIC), each of which may have different cell sizes, the cell size is a programmable value in units of 32 bytes.

(72) There may be two sets of VCs with each SQ **330** assigned to one set. A maximum frame size worth of space can be reserved for each VC, and each VC set can have a different maximum frame size. The remainder of the link partner input buffer is shared dynamic space usable by any SQ/VC, subject to per VC and buffer class limits.

(73) The size that comes with the request represents the size of the packet on the wire which includes a 4-byte FCS. This gets converted to an internal 2-byte FCS at the link partner before

writing the packet to the link partner input buffer so the crediting needs to account for this difference, which can be a factor at the boundary of the cell size. For instance, for a 96 byte cell, a size that is 97 or 98 will take a single cell. In order to know when this happens, the request includes a correction term which is calculated as: $\text{req_len_correct} = (\text{byte_len} \% 16) == 1 \text{ or } 2$.

(74) Further validation of this term is required to convert it to whatever the cell size boundary may be. It will be valid when the length just exceeds the cell size. With this, the validated fen_correct term can be determined by: $\text{len_correct} = (((16 - \text{byte size}) \% (2 * 32 - \text{byte cell size})) == 1) \& \text{req_len_correct}$

(75) An example of how these values work for a few cell and packet sizes is illustrated in the table below:

(76) Length Correct Calculation

(77) TABLE-US-00002

len_correct	Size	Req	Size	Cell	Size	Modulo	Credit	(bytes)	len_correct																											
(16B units)	(32B units)	result	len_correct	Taken	64	0	4	2	0	2	65	1	5	2	1	1	2	66	1	5	2	1	1	2	67	0	5									
2	1	0	3	96	0	6	3	0	0	3	97	1	7	3	1	1	3	98	0	7	3	1	0	4	128	0	8	4	0	0	4	129	1	9	4	1
4	130	1	9	4	1	1	4	131	0	9	4	1	0	5																						

(78) The size that comes with the request uses 8-byte units and the link partner input buffer cell size is a multiple of 32 bytes ($32 * y$ where $y = \text{cell size from CSR}$). First, the 8-byte size is converted to a 16-byte size ($\text{ROUNDUP}((8 - \text{byte size}) / 2)$). Also, the cell size is converted to 16 byte units ($2 * y$). Mathematically, the number of cells a request will use can be calculated by:

$\text{ROUNDNDN}(((16 - \text{byte size}) + 2 * y - 1 - \text{len_correct}) / (2 * y)) = \# \text{ of cells}$

(79) While a divide operation is possible in hardware, due to timing reasons, a divide operation cannot be done in the critical path of the arbitration. Instead, an alternate credit management is used. That is, credits are maintained in units of 32 bytes. When a request wins arbitration, the number of credits taken is adjusted by the maximum error term ($2 * y - 1$) using the calculation: $\text{ROUNDNDN}(((16 - \text{byte size}) + 2 * y - 1) / 2) = \text{Maximum 32 byte credits needed}$. Because this calculation overestimates the credit required for the packet, on the following clock, a modulo operation ($X = (16 - \text{byte size}) \text{ MOD } 2 * y$, $y = 32 - \text{byte cell size from CSR}$) can be performed to determine what the actual remainder is. This value along with the len_correct term are used to adjust the credit counter. The formula used to create the adjustment value (adj_val) for X is: If ($X == 0$) $\text{adj_val} = y - 1$ else if ($X == 1$ and fen_correct) $\text{adj_val} = y$ else $\text{adj_val} = \text{ROUNDNDN}((X - 1) / 2)$

(80) The table below illustrates a request credit example for 96 byte cells showing the values used across several packet lengths for the 96 byte cells of the switch input buffer ($y = 3$).

(81) Request Credit Example For 96-byte Cells

(82) TABLE-US-00003

Packet	Size	Packet	Size	Credit	Modulo	Corrected	(bytes)	(16-byte units)																						
Taken	Result	len_correct	adj_val	Credit	Taken	48	3	4	3	0	1	3	64	4	4	4	0	1	3	80	5	5	5	0	2	3	96	6	5	0
0	2	3	97	7	6	1	1	3	3	98	7	6	1	1	3	3	99	7	6	1	0	0	6	128	8	6	2	0	0	6

(83) If a request is filtered before being forwarded to the link partner input buffer, the output buffer logic returns the SQ and VC so they can be used to return the credits to the appropriate credit counters. No size is required since the packet size is always the same, the length of a reduction frame (69-byte or 16-byte size=5).

(84) The local (master) side of the link maintains a count of the number of packets sent from each VC across both sets (8 total), a count of amount of packet (in 3 2-byte quantities) sent to each VC (4), and a count of the amount of packet (in 32-byte quantities) sent for each buffer class (8). The link partner (slave) side of the link maintains the same set of counts with them being sent over the link periodically. The difference between the master and slave counts is a count of the number of packets in the link partner input buffer from each VC across both sets and a count of the amount of space (in 32-byte quantities) currently occupied by each VC and each buffer class. A count is also maintained of the total amount of space used across all packets. A summary of the counters is as follows: $\text{master_vcx_cnt}[4] / \text{slave_vcx_cnt}[4]$ —master and slave counts of the number of packets sent to each VC in set X; $\text{master_vcy_cnt}[4] / \text{slave_vcy_cnt}[4]$ —master and slave counts of the

number of packets sent to each VC in set Y; master_bc_cnt[8]/slave_bc_cnt[8]—master and slave counts of the amount of space occupied by each buffer class in units of 32-bytes; master_vc_cnt[4]/slave_vc_cnt[4]—master and slave counts of the amount of space occupied by each VC in units of 32-bytes; master-tot-cnt/slave-tot-cnt—master and slave counts of the total amount of space occupied in units of 32-bytes.

(85) All counters are set to zero on a warm reset. They are also forced to zero when the link is in the draining state or when the DBG_RESET CSR bit to clear their state is set. The output buffer filter will steer a reduction packet to something other than the path to the link partner input buffer. In this case, a signal can be returned along with the SQ and VC of the packet. Again, the length is not required as the size of these packets is fixed. This information is used to adjust the appropriate master credit counts.

(86) A request is allowed to participate in arbitration if either its VC count is 0 (indicating its one statically assigned slot is available) or there is space for a max sized frame in the dynamic space (subject to the targeted buffer class and VC limits). There can be a single programmable value for max frame size which is used across all VCs and SQs. The request validation for input buffer space can be addressed using credit-based flow control.

(87) Credit-based flow control can be used to divide a dynamic space in two ways, each independent of each other: first, based on a limit of how much dynamic space each of the four VCs can take; and second, based on a limit to how much dynamic space each of the eight buffer classes can take. In both cases, the limits are set as a percentage of the available space. For a given packet, space should be made available in both its targeted VC and buffer class. For instance, if each space has its limit set to 50%, if one is active, it has access to 50% of the buffer space, with two active, each space gets 37.5% $((50+50*0.5)/2)$, with three active, each space gets 29.2% $((75+25*0.5)/3)$, and so on. Also, the total space used by those spaces that are active can be limited to the given total (50%, 75%, 87.5%). Accordingly, the space allocated to each varies dynamically by how many are currently active. When an additional one goes active it causes others that are active to give up some of their space which is then taken by the new one.

(88) Like the division function discussed above, this function is implemented as a lookup table. For the VC space in this example, there are 16 entries with each entry specifying the space available to each VC along with the total space available across all VCs. For the buffer classes, there may be 256 entries with each entry specifying the space available to each buffer class along with the total space available across all buffer classes. Space for each is expressed in 2048-byte units. The depth of each table is sufficient to cover all combinations of active members (VCs or buffer classes), with each being able to have an independent setting for their percentages. With this, it is up to software to program the values in the table according to how much total dynamic space there is and what percentage each is allowed to take across all possible combinations. As more become active, each is allowed less space and the total available increases. Requests for spaces that are above this limit, or in total above the total limit, are not allowed to take more dynamic space.

(89) A VC or buffer class is considered active either if it has a request in an age queue, or if it has outstanding credits for link partner input buffer space. As an example, consider there are only 4 spaces (16 entry table) with percentages set as SPACE0(50%), SPACE1(40%), SPACE2(30%), SPACE3(10%), and a total dynamic space of 16 KB. This results in the values, in quantities of 16-bytes presented in the buffer space example table below.

(90) Buffer Space Example

(91) TABLE-US-00004

Index	SPACE3	SPACE2	SPACE1	SPACE0	Total	0	N/A	N/A	N/A	N/A	N/A
1	N/A	N/A	N/A	512	512	2	N/A	N/A	410	N/A	410
3	N/A	N/A	N/A	319	398	717	4	N/A	307	N/A	N/A
5	N/A	250	N/A	416	666	6	N/A	255	339	N/A	594
7	N/A	202	270	337	809	8	102	N/A	N/A	N/A	102
9	94	N/A	N/A	469	563	10	94	N/A	377	N/A	471
11	75	N/A	299	374	748	12	95	284	N/A	N/A	379
13	78	234	N/A	389	701	14	80	239	319	N/A	638
15	79	236	315	394	1024						

(92) As an example, the values in the row for index 7 are calculated as: Total % = $0.5 + (1 - 0.5) * 0.4 +$

$(1-0.5-(1-0.5)*0.4)*0.3=0.79$; $SPACE0=(0.5/(0.5+0.4+0.3))*0.79*1024=337$; $SPACE1=(0.4/(0.5+0.4+0.3))*0.79*1024=270$; $SPACE2=(0.3/(0.5+0.4+0.3))*0.79*1024=202$;
Total=337+270+202=809

(93) As noted above, and referring back to FIG. 2, switches, such as switch **202** may be used to create a switch fabric, where the switch ports **220** may be configured to operate as either edge ports or fabric ports. As also noted above, switch **202** can support various network topologies including but not limited to, e.g., dragonfly and fat-tree topologies. Networks can be thought of as comprising one or more slices, each having the same overall topology, although slices may differ with respect to how each is populated. Nodes are connected to one or more ports on each slice. When a network has multiple slices, and a node is connected to more than one slice, the node is assumed to be connected at the same location in each slice.

(94) Routing in the switch fabric may be controlled by a fabric routing function (FRF) implemented in switch **202**. An example FRF component **400** is illustrated in FIGS. 4A and 4B. It should be understood that a separate instance of FRF component **400** may be implemented within input logic for each port of switch **202**. Routing decisions made by FR component **400** can be applied to those frames that are not already part of an established flow. It should be noted that FRF component **400** does not necessarily know whether or not a particular frame is associated with a flow, but rather makes an independent forwarding decision for each frame presented at an input port. FRF component **400** may comprise filters, tables, circuitry, and/or logic such as selection circuitry/logic to effectuate routing of data throughout a switch fabric as described herein. As illustrated, FRF component **400** includes at least: a minimal ports selection component **402** (which includes a minimal tables component **402A**), various ports filters (permitted ports filters, operational ports filters, busy ports filters); a preferred ports discrimination component **402B**; pseudo-random down selection components/logic **402C**; exception tables **404** (including an exception list table **404A**); operational ports component **406** that includes a global fault table **406A**; and a routing algorithm table **408**. As illustrated in FIG. 4B, FRF component **400** may further comprise: a non-minimal ports selection component **410** that includes local non-minimal selection component **410A** and global non-minimal selection component **410B**; and output logic component **412** (which is part of a switch's output control block), which includes an adaptive selection component or logic **412A**. FRF component **400** includes other components, and are described herein.

(95) In particular, FRF component **400** determines a preferred port with preferred ports discriminator **402B** to forward each frame presented at the input port based on: a received frame's destination fabric address (DFA); the frame's current routing state (where the frame is along its path, and the path(s) it took to reach its current routing state); the switch fabric routing algorithm and configuration; and load metrics associated with the output port (the aforementioned preferred port to which the frame is to be forwarded) using busy ports filters.

(96) FRF component **400** may include a routing algorithm table **408** that may be embodied as a software configurable table that determines valid choices based on the frame's current routing state. Valid choices are decisions such as whether a local minimal, global minimal, local non-minimal, or global non-minimal path is allowed to be chosen for the frame's next hop. The routing state includes information such as the VC the frame was received on, and whether it is in the source, the destination, or an intermediate group. The routing algorithm table **408**, along with the adaptive selection function or logic **412A** (described below), also determines the VC to be used for the frame's next hop.

(97) Frame routing with unicast DFAs will be described as an example. However, it should be noted that the DFA of the routing request can either be in unicast or multicast format. The unicast format can include a 9-bit global ID field (global_id), a 5-bit switch ID field (switch_id), and a 6-bit endpoint ID field (endpoint_id). The global ID can uniquely identify a group within the network. Specifically, it identifies the final group to which the frame must be delivered. The switch

ID uniquely identifies a switch within the group identified by global ID. The endpoint ID field, together with the global ID and switch ID identify the endpoint, connected to the edge of the network fabric, to which the frame is to be delivered. This field is mapped to a port or set of ports on the switch identified by global ID and switch ID.

(98) The multicast format includes a 13-bit multicast ID field (multicast_id). This field is mapped by FRF component **400** to a set of ports on the current switch to which the frame is to be forwarded.

(99) From this information, FRF component **400** determines an updated routing state for the frame, which is then carried within the frame. For example, to effectuate routing in a dragonfly topology, a frame's current state may be gleaned from the frame's VC (discussed above). Based on algorithmic switch fabric routing rules specified for the switch fabric (the selection of which is described below), FRF component **400** determines a particular VC to be used for the frame's next hop to avoid any deadlocks. Additional routing state information can be provided depending on where the frame is along its path, e.g., whether the frame is in its source group, in an intermediate group, or in its destination group. It should be noted that FRF component **400** performs port filtering (described in greater detail below) using permitted ports filter, operational ports filter, busy ports filters, etc. to determine if a preferred port to which a frame is to be forwarded is currently faulty, busy, absent, etc.

(100) Switch **202** distributes load information between switches. The FRF component **400** receives the load measurement of and from its associated output port. The FRF component **400** receives summary load information from its associated input port for a neighboring switch. Each FRF component **400** exchanges load information with all other FRF instances within the same switch. FRF component **400** provides summary load information to its associated output port for delivery to a neighboring switch. Through the load distribution mechanism, each FRF component **400** learns the load measured at each output port of its switch. As well, each FRF learns the summary load information for all neighboring switches.

(101) It should be noted that FRF component **400** can support frame multicasting. When a multicast DFA is received, FRF component **400** determines a set of ports to which the frame associated with the multicast DFA should be forwarded. The set of ports can be determined by accessing a lookup table that maps software-configured multicast fabric addresses to output ports. This helps avoid problems associated with duplicate multicast frame copies.

(102) FIG. 5 illustrates an example route selection process involving down-selection of candidate ports and adaptive route selection based on load. FRF component **400** considers three categories of candidate ports to which a frame may be forwarded: preferred minimal path candidate ports **502**; non-preferred minimal path candidate ports **504**; and non-minimal path candidate ports **506**. Depending on where a frame may be along its path, non-minimal path candidate ports are either global, non-minimal path candidate ports or local non-minimal path candidate ports.

(103) Filtering may be applied to the three categories of candidate ports, e.g., operational port filtering, useable port filtering, and busy port filtering. Port filtering as applied herein can be used to reduce the set of valid ports considered as path candidate ports by identifying and removing absent and/or faulty ports and/or busy ports from consideration.

(104) Operational port filtering (or non-operational port filtering) can refer to the removal of non-operational ports from sets of ports being considered as candidates for routing, e.g., preferred minimal path candidate ports **502**, non-preferred minimal path candidate ports **504**, and non-minimal path candidate ports **506**. That is, switch **202** may identify certain ports as being non-operational. These non-operational ports may be reported in a non-operational port mask. It should be noted that in some embodiments, software can force additional ports of switch **202** to be considered as non-operational using a non-operational port CSR, e.g., when a port(s) is to become disconnected as a result of planned maintenance.

(105) Usable (or unusable) port filtering may involve filtering out candidate ports for consideration

that would normally have been acceptable, but due to faults within network **100**, for example, the candidate ports have become unacceptable/unusable for reaching one or more destination switches, destination groups (of switches), etc., but remain acceptable or usable for reaching one or more other destination switches. In some embodiments, global fault table **406A** can be used to block global minimal path port candidates and global non-minimal path port candidates depending on a destination group of the frame being routed. For example, candidate ports that lead to an intermediate group (of switches) without connectivity to a particular destination group (of switches) can be excluded from consideration when routing frames to that destination group, although the same candidate ports may not necessarily be blocked out for other destination groups. The global fault table **406A** can be determined or indexed by the `global_id` field of the frame's DFA.

(106) In some embodiments, an exception list maintained by exception list table **404A** may be used to conditionally exclude port candidates depending on the destination group or switch to which the frame is being routed. It should be noted that exception list table **440A** may be used to identify preferred global minimum path ports. Accordingly, use of the exception list table **404A** to exclude port candidates is done when it is not being used to identify preferred global minimum path ports.

(107) It should be noted that knowledge regarding which ports are busy in a neighboring switch can be used to determine if the ports that connect to a neighboring switch are poor candidates to receive a frame based on whether the neighboring switch will subsequently need to forward the frame to a port that is already busy. For example, when considering candidate ports for use in global minimal routing, the ports connected to a neighboring switch are poor candidates if the neighboring switch's global ports that connect to the frame's destination group are all busy. Similarly, when in the destination group and considering candidate ports for use in local non-minimal routing, the ports connected to a neighboring switch are poor candidates if the neighboring switch's local ports that connect to the frame's destination switch are all busy.

(108) Accordingly, busy port filtering can be performed by FRF component **400** by using busy port masks to remove heavily loaded ports from being considered candidate ports. It should be noted that in some embodiments, heavily loaded ports are removed from consideration when other candidate ports that are not heavily loaded exist. Otherwise, when non-heavily loaded ports do not exist, busy port filtering will not remove the heavily loaded ports from consideration. FRF component **400** maintains four masks of busy ports, that is, ports whose load exceeds a software-defined threshold: local switch busy port mask; global non-minimal busy global port mask; global non-minimal busy local port mask; remote switch busy port mask. The remote busy port mask and other load information is communicated between switches to populate the remote switch.

(109) A local switch busy port mask can be applied to minimal path candidate ports as well as to local non-minimal path candidate ports. The FRF generates a 64-bit `Is_busy_port_mask` by comparing each port's `local_load` to a software defined threshold. Ports with loads higher than this threshold are marked as busy in this mask.

(110) A global non-minimal busy port global mask can be applied to global ports of global non-minimal path candidate ports. The FRF generates a 64-bit `gnmbgp_mask` by comparing each port's `gnmgrp_load` to a software defined threshold. Ports with loads higher than this threshold are marked as busy in this mask.

(111) A global non-minimal busy local port mask can be applied to local ports of global non-minimal path candidate ports. The FRF generates a 64-bit `gnmblp_mask` by comparing each port's `gnmlp_load` to a software defined threshold. Ports with loads higher than this threshold are marked as busy in this mask.

(112) A destination group dependent busy-port mask, obtained from a remote switch busy global port table can be applied to global minimal path candidate ports. Correspondingly, when a frame is being routed in its destination group, a destination switch-dependent busy-port mask, obtained from a remote switch busy local port table can be applied to local non-minimal path candidate

ports.

(113) Upon applying the aforementioned filtering or down-selection stage, a set of surviving path candidate ports **508** can result. That is, a reduced number of candidate ports can be determined after removing non-operational and unusable heavily loaded port candidates a set of path candidate ports remains. In some embodiments, a pseudo-random selection process is used to further reduce the number of surviving path candidate ports **508** to a determined threshold number of ports associated with each category of candidate ports (preferred minimal path candidate ports, non-preferred minimal path candidate ports, and non-minimal path candidate ports). In some embodiments, that threshold number of candidate ports may be four candidate ports per category. If the minimum threshold number of candidate ports is not met, no candidate ports from that category are removed from consideration.

(114) In some embodiments, this pseudo-random selection (or down selection) of candidate ports can be weighted. As will be described below, this weighting can be applied to counteract a potential bias in the distribution of globally non-minimal traffic among a group's global links. That is, weights, e.g., weights between 0 and 15, can be assigned to each port per the CSR configuration. This weighting can be used to influence the probability with which individual candidate ports are chosen such that higher-weighted ports have a great chance of being chosen. For example, a weight of 15 results in a port having 15 times greater likelihood of being selected in the pseudo-random selection process. In some embodiments, candidate ports may be filtered into four groups (GW1, GW2, GW4, GW8) based on their assigned weights, where a candidate port can belong to multiple groups depending on the assigned weight (e.g., a candidate port with weight 1 belongs only a one weight group, while a candidate port with weight 5 belongs to two groups (GW1 and GW4, i.e., $1+4=5$), and a candidate port with weight 15 belongs to all four groups (GW1, GW2, GW4, GW8, i.e., $1+2+4+8=15$). The number of candidate ports in each group can be determined ($nW1$, $nW2$, $nW4$, $nW8$), and pseudo-random selection is applied to each group to select one candidate port from each group ($cW1$, $cW2$, $cW4$, $cW8$). The weight of each group can be computed, along with their total weight: $wW1=nW1$; $wW2=2*nW2$; $wW4=4*nW4$; $wW8=8*nws$; $wtotal=wW1+wW2+wW4+wW8$. A fifth pseudo-random selection can be performed to choose a number j in the range $0 \dots Wtotal-1$. One of the candidates, $cW1$, $cW2$, $cW4$, $cW8$ is chosen as the down selected candidate port based on the value of j as follows: If $j < wW1$, choose $cW1$; else If $j < wW1 + wW2$, choose $cW2$; else If $j < wW1 + wW2 + wW4$, choose $cW4$; else choose $cW8$.

(115) It should be noted that the same candidate port may be chosen by more than one of the instances/iterations of the weighted pseudo-random down selection logic **410C**, in effect reducing the number of candidates ports that are chosen. The probability of the same candidate port being chosen by more than one of the $mgnm=4$ global non-minimal, weighted pseudo-random selectors decreases with increasing numbers of candidates ports to choose from. In the context of a dragonfly topology, for example, and a network with full global bandwidth, at an edge port in the source group there are potentially about 48 possible global non-minimal candidates ports: 16 global ports and 32 local ports. If a local hop has been taken, the next hop is a global hope, thereby reducing the number of candidate ports to about 16. However, if the network is tapered such that it only supports one quarter of full global bandwidth, there may only be 4 global candidates to choose from after a local hop has been taken.

(116) FRF component **400** can use received remote switch busy port masks to generate the aforementioned Remote Switch Busy Global Port Table of busy port masks that identify ports connected to neighboring switches that should be avoided as global minimal path candidates based on the destination group to which the frame is being routed. Similarly, the received remote switch busy port masks can also be used to generate the aforementioned remote switch busy local port table of busy port masks that identify ports connected to neighboring switches that should be avoided as local non-minimal path candidates, when routing in the destination group, based on the destination switch to which the frame is being routed.

(117) The `rs_busy_port_masks` are used in assessing the suitability of neighboring switches (whether ports of neighboring switches are busy or quiet) for reaching specific destination groups via global minimal paths and specific destination switches via local nonminimal paths or via local minimal paths. Each FRF instance corresponding to a local port or a global port can be configured to generate a 64-bit `rs_busy_port_mask`. The generated mask is delivered to the partner switch connected to that port. Similarly, the partner switch can also generate and return an `rs_busy_port_mask`.

(118) Each FRF instance communicates the `rs_busy_port_mask`, that it received from its partner switch, to the FRF instances in all of the switch's other ports using the port status ring. The port status ring communicates load and status information between all the FRF instances in the switch. Each FRF instance captures all `rs_busy_port_masks` such that all FRF instances learn the remote busy port status being provided by all neighboring switches. Each FRF instance uses the `rs_busy_port_masks` that it receives to generate the Remote Switch Busy Global Port (RSBGP) Table and Remote Switch Busy Local Port (RSBLP) Table.

(119) Generation of the `rs_busy_port_mask` is a two-step process. The first step is to compare each port's `local_load` to a software configurable threshold generating an intermediate mask of all ports that are individually busy. Ports that are classified as non-operational are also recorded as busy in the intermediate mask. The second step takes link bundling into account such that a port is only marked as busy in the `rs_busy_port_mask` if it and all other ports, that are part of the same bundle, are marked as busy in the intermediate mask. Either all ports that are members of the same bundle are marked as busy in the `rs_busy_port_mask`, or none are. Global ports that are part of the same bundle all connect to the same remote group. Local ports that are part of the same bundle all connect to the same remote switch within the current group.

(120) As the `rs_busy_port_masks` are used to determine whether the switch that generated the mask is a good candidate for routing a frame to another group or to another switch in the current group, bundling is used to provide a consistent view of the generating switch's suitability when the busy status across its links, that connect to the destination group or to the destination switch in the current group, is inconsistent. The rationale for the treatment of bundling described here is that the switch generating the `rs_busy_port_mask` remains a candidate for reaching the destination group, or the destination switch in the current group as long as it has at least one link to the destination group or switch that is not busy; adaptive routing at the switch that generated the `rs_busy_port_mask` should direct the frame to the non-busy link.

(121) Ports must be included in either the bundled ports mask CSR or the unbundled ports mask CSR (both of which are part of the static description of the wiring) in order for them to be marked as busy in the `rs_busy_port_mask`. The bundled port masks are scanned to identify the bundles and the ports that they contain. In addition, the unbundled port mask is consulted to identify any other ports, that are not members of a bundle, but whose busy status should also be included in the generated `rs_busy_port_mask`.

(122) A different software-defined threshold is used in computing the `rs_busy_port_mask` because of the larger latency involved in communicating and processing the `rs_busy_port_mask` and in delivering a frame subject to this mask to the remote switch that generated the mask. Because of the larger latency, it may be useful to require a port to be more loaded before it is considered to be so busy that it is not a good candidate for receiving additional frames from a remote switch. A busy remote port should be sufficiently loaded such that it remains loaded throughout the time it would take to receive a frame that has been subject to the mask.

(123) The aforementioned RSBGP table stores busy port masks indexed by destination group (`global_id`). It is populated dynamically using information contained in received `rs_busy_port_masks` together with topological information about the system stored in configuration state. Again, the RSBGP table is used in the evaluation of global minimal paths consisting of a hop to a neighboring switch which has a global port connected to the destination group to filter out

ports of the current switch which are poor choices for use in reaching the destination group because the corresponding global port or ports, of the neighboring switch reached by the filtered-out ports of the current switch, are too heavily loaded.

(124) The RSBLP table stores busy port masks indexed by destination switch (`switch_id`), and again, can be used in the evaluation of local non-minimal paths consisting of a local hop to a neighboring switch followed by another local hop to the destination switch. It is populated dynamically using information contained in received `rs_busy_port_masks` together with topological information about the system stored in configuration state. For topologies, such as fat-tree, where a local minimal path can consist of a local hop to a neighboring switch followed by another local hop to the destination switch, the RSBLP Table can also be used in the evaluation of local minimal paths. The RSBLP table is used to filter out ports on the current switch which are poor choices for use in indirectly reaching the destination switch because the neighboring switch's port or ports that connect to the destination switch are too heavily loaded.

(125) It should be noted that the RSBGP table and the RSBLP table are never both accessed for the same routing request. The former is accessed when routing a frame that is not in the destination group, and the latter is accessed only when routing a frame that is in the destination group. Therefore, both can be implemented within the same memory, termed the Remote Switch Busy Port Table.

(126) A Busy Ports Filter is used to remove candidate minimal path ports that are poor choice because they lead to heavily loaded paths. Separate copies of the Busy Ports Filter are applied independently to the preferred and to the non-preferred minimal path ports. If global minimal routing is enabled a first stage of this filter removes ports marked as busy in the indexed entry of the Remote Switch Busy Global Port Table. If in the destination group and local minimal routing is enabled, an optional first stage of this filter removes ports marked as busy in the indexed entry of either the Remote Switch Busy Local Port Table or the Remote Switch Busy Global Port Table, with the table that is used being based on CSR configuration. The first stage of this filter is not applicable and not used with the Dragonfly topology when in the destination group; however, it can be applicable to other topologies including fat tree. A second stage of this filter removes ports marked as busy in the aforementioned `Is_busy_port_mask`. If no candidates remain after the second filter stage, the second filter stage is disabled such that the only candidates ports that remain, if any, will be seen to have a high load in the Adaptive selection stage.

(127) A Local Non-Minimal Busy Port Filter is used in the destination group to remove local non-minimal candidate ports that are poor choice because they lead to heavily loaded paths. A first stage of this filter removes ports marked as busy in the indexed entry of either the Remote Switch Busy Local Port Table or the Remote Switch Busy Global Port Table, with the table that is used being based on CSR configuration. A second stage of this filter removes ports marked as busy in the aforementioned `Is_busy_port_mask`. If no candidates remain after the second filter stage, the second filter stage is disabled such that the only candidates ports that remain, if any, will be seen to have a high load in the Adaptive selection stage.

(128) A Global Non-Minimal Busy Port Filter is used when global non-minimal routing is enabled to remove candidate global non-minimal path ports that are poor choice because they lead to heavily loaded paths. Specifically, global ports marked as busy in the aforementioned `gnmbgp_mask` are removed and local ports marked as busy in the aforementioned `gnmblp_mask` are removed. If no candidates remain after application of these masks, both masks are disabled such that the only candidates ports that remain, if any, will be seen to have a high load in the Adaptive selection stage.

(129) Provided that there is at least one valid candidate port, the various busy port filters (Busy Ports Filter, Local Non-Minimal (LN) Busy Port Filter, Global Non-Minimal (GN) Busy Port Filter) may not be allowed to collectively block all candidate ports. If there are viable candidate port choices, they are allowed, despite being “poor” choices, if there are no better choices.

Otherwise, an empty route response will be generated for the routing request when routes are actually available.

(130) To prevent an incorrect empty route response from being generated, the first stage of the preferred and non-preferred minimal path busy port filters (FIG. 4A) and the first stage of the local non-minimal path busy port filter (FIG. 4B) are all disabled if the following conditions are all true: No candidates survive the first stage of the preferred minimal path busy port filter (Busy Ports Filter); No candidates survive the first stage of the non-preferred minimal path busy port filter (Busy Ports Filter); No candidates survive the first stage of the local non-minimal busy port filter (Local Non-Minimal (LN) Busy Port Filter); and No candidates survive the global non-minimal busy port filter (Global NonMinimal (GN) Busy Port Filter).

(131) It should be noted that there will be no minimal path candidate ports if minimal routing is disabled (Permitted Ports Filter). There will be local non-minimal path candidate ports only if local non-minimal routing is enabled (Candidate Local Non-Minimal Path Ports). There will be global non-minimal path candidates only if global non-minimal routing is enabled (Candidate Global Non-Minimal Path Ports). Local and global non-minimal routing are generally not both enabled simultaneously. When the first stage of the preferred and non-preferred minimal path busy port filters and the local non-minimal busy port filter are disabled due to the conditions described above, the only candidate ports that will be seen at an adaptive selection stage (described below) will be poor candidates because they will all be ports that lead to other switches whose ports (that connect to the destination group or to the destination switch) are heavily loaded. This is because these are the only candidate ports that were being blocked by the filters that are being disabled and, without these filters being disabled, there are no other candidates.

(132) The adaptive selection stage will choose between these remaining/surviving candidate ports, which are all poor, based on their biased local loads, although their local loads will not necessarily reflect the reason why they are poor. Their poor character can be the result of high downstream load on certain ports of the other switches reached by these candidate ports. It is because the adaptive selection stage may not be able to see how poor these candidates are that the coordination between the different busy port filters, described herein may be used. If each busy port filter decides independently whether or not to disable its RSBGP Table and RSBLP Table-based filters, situations such as the following could occur. The non-preferred minimal path busy port filter might produce one or more candidates, which are not poor, without any of its filter stages disabled. The preferred minimal path busy port filter might only be able to produce one or more candidate ports by disabling both of its filter stages. Thus, all of the candidate ports it is able to produce are poor. At the adaptive select stage, the down-selected, not poor, non-preferred minimal path candidate ports, are compared against the down-selected, poor, preferred minimal path candidate ports. However, the adaptive selection stage lacks visibility into how poor the preferred minimal path candidates are, so it may select a poor preferred minimal path candidate over a not poor non-preferred minimal path candidate.

(133) An alternative to the busy port filter coordination mechanism described herein, would be for all of the busy port filters to act independently, but for the minimal path and the local non-minimal path busy port filters to each forward a signal through to the adaptive select stage to indicate if their respective candidate ports are poor choices due to busy ports at downstream switches. If they are, the adaptive selection function may de-prioritize their candidate ports in favor of other ports. The result would be the same as is achieved by the coordination, described herein, between the various busy port filters.

(134) As illustrated in FIG. 5, load-based adaptive selection can be performed on the surviving path candidate ports **510** that remain after the pseudo-random selection process is performed by FRF component **400**. The adaptive selection stage will result in a single, least loaded candidate port **512** to which a frame can be routed, where the current load present on the candidate ports surviving pseudo-random down selection (surviving path candidate ports **508**) are compared to determine the

least loaded candidate port among this remaining set of candidate ports.

(135) In some embodiments, preferred minimal path candidate ports are preferentially selected over non-preferred minimal path candidate ports, and minimal path candidate ports are to be preferentially selected over non-minimal path candidate ports. To accomplish this preferential selection, a bias value can be added to each candidate port's load before the adaptive selection comparison is performed. The bias value used can be configured using CSRs, and can vary depending on the type of path to which it is being applied (i.e., non-preferred minimal, preferred minimal, an non-minimal), the traffic class of the frame being routed, and where the frame is along its path. For example, frames belonging to a low-latency traffic class can be more strongly biased towards minimal paths versus frames in other traffic classes to have a greater likelihood of achieving/comporting with low-latency requirements or needs. Frames may also be increasingly biased towards minimal path routes the closer the frames are to their destination.

(136) In particular, load values represent the busyness of switch **202**'s ports and are used in evaluating the load-based port masks and in comparing candidate ports during the adaptive route selection process. Load-based port masks are used in the busy port filters to remove ports that are poor candidates, based on current load, from the set of candidates being considered. There are a number of different types of load values used within the switch and some are communicated to neighboring switch devices. These load values are described in the following sections, noting that FIG. **6** illustrates how the load values are measured, communicated, and used.

(137) Several load metrics are computed and used in determining which port to route a frame when there is more than one port to which the frame can be routed. The load metrics are also used in generating busy-port masks, which as described above, are used to remove heavily loaded ports from consideration.

(138) There are five load metrics described herein: local load; group load; global non-minimal global port load; mean global load; and global non-minimal local port load.

(139) Regarding local load, the load of each of a switch's output ports (e.g., output ports **220c** of switch **202**) is continuously being evaluated and provided to the corresponding FRF instance as the 8-bit value `local_load`. Larger values represent higher load. The current load present at each output port is measured by the output control age queue block. The output port load is provided by each age queue instance to the FRF instance (of FRF component **400**) that is associated with the input side of the same port. The load value provided to the FRF is an 8-bit value termed the `local_load`. The age queue determines the local load based on a combination of the amount of traffic enqueued, waiting to go out that port and the amount of traffic enqueued at the opposite side of the link in the link-partner switch device's input buffer. Each FRF instance distributes the `local_load` value it receives from its associated age queue instance to all other FRF instances. In this way, each FRF instance learns the current `local_load` of every output port.

(140) When the port loads of candidate ports are being compared to determine the best port to route a frame to, it is the port's `local_load` value that is used for ports being considered for minimal path routing, and for ports being considered for local non-minimal path routing.

(141) Group load is a measure of how suitable a dragonfly group is for use as the intermediate group in a global non-minimal path. The 8-bit group load value is not computed by a switch, such as switch **202**, but is software-configurable. Software might use a measure of the network injection load present across the input side of the group's edge ports in deriving the `group_load` value, or might determine the `group_load` value based on a policy of discouraging use of certain groups as non-minimal intermediate groups, perhaps based on the jobs or services that are running in the groups. That is, the `group_load` value is intended to be representative of the amount of local traffic within a group.

(142) A network management stack sets the `group_load` value by periodically writing to a CSR. The software-configured group load value is communicated across global links. FRF instances associated with global links forward the `group_load` value that they receive from their link-partner

in the group at the opposite side of the link, to all other FRF instances in the switch. In this way, each FRF instance learns the group load values of the groups at the opposite end of each of the global links terminated by the switch.

(143) In terms of global non-minimal global port load (`gnmgrp_load`), this load is a metric used in assessing a global port's suitability for use in directing a frame to the intermediate group reached by the global link connected to the global port. The `gnmgrp_load` is nominally equal to the maximum of the global port's `local_load` and the `group_load` value being received from the group reached by the global link. However, through CSR configuration, the `group_load` component can be excluded.

(144) When the port loads of candidate ports are being compared to determine the best choice port to route a frame to, it is the port's `gnmgrp_load` value that is used for global ports being considered for global non-minimal path routing.

(145) Mean global load (`mean_global_load`) is intended for use in assessing a switch's suitability for use in reaching any intermediate group that is directly connected to that switch. The `mean_global_load` value is an 8-bit value equal to the arithmetic mean of the `gnmgrp_load` values of all of the switch's global ports. Ports that are classified as non-operational by either hardware or software are excluded from the calculation. The ports to include in the `mean_global_load` computation are determined from through CSR configuration.

(146) For any port whose load is included in the `mean_global_load` computation, if a `group_load` value is not being received for that port, either because link partner data is not being received for that port or because the link partner data that is being received is not global link data, the contribution of that port to the `mean_global_load` is based solely on that port's local load. The computed `mean_global_load` value is communicated across local links. FRF instances associated with local links forward the `mean_global_load` value, that they receive from their link-partner in the local switch at the opposite side of the link, to all other FRF instances in a switch. In this way, each FRF instance learns the `mean_global_load` values of the local switches at the opposite end of each of the local links terminated by the switch, and can use these values in global non-minimal path selection.

(147) Global non-minimal local port load metrics are computed by each FRF instance. The global non-minimal local port load (`gnmlp_load`) is a metric for use in assessing a local port's suitability for use in directing a frame to an intermediate group of a global non-minimal path. The global non-minimal local port load takes into account the load on the local port as well as the suitability of the local group switch, to which the port connects, for use in reaching an intermediate group. A port's `gnmlp_load` value is equal to the maximum of the port's `local_load` and the `mean_global_load` reported by the port's partner switch. Through software configuration it is possible to remove the `mean_global_load` component such that a port's `gnmlp_load` becomes simply equal to its `local_load`.

(148) The `gnmlp_load` value is an 8-bit value. Each FRF instance computes the `gnmlp_load` value for all of its switch's ports at which link partner data for a local link is being received. If load status information is not being received from a port's partner switch, the `gnmlp_load` value for that port is set equal to the port's `local_load`.

(149) FIG. 6 illustrates example load measurements and how load measurement may be exchanged between switches in a group. FIG. 6 illustrates a group of switches, e.g., group 1, comprising switches **602**, **604**, and **606**, each of which may be embodiments of switch **202** (FIG. 2). Group load values, as noted above, can be exchanged across global links, and as shown, `group_load` values are transmitted from/received by each of switches **602**, **604**, and **606** from other groups/switches in the switch fabric. Within group 1, switches **602**, **604**, and **606** exchange `mean_global_load` values. Switch **602** also is shown as computed `gnmlp_load` values at the output side of local ports, based on `local_load` at that port, and `mean_global_load` reported by the link partner. Further still, `local_load` values at the output of all ports are computed by switch **602**. It should be noted that `mean_global_load` and `group_load` values are sent across links between

connected switches symmetrically from each switch to the other.

(150) The fabric routing process described in the preceding subsections is performed for every frame received. Switch determines whether to perform packet-by-packet adaptive routing (using the recommended port choice) or flow based adaptive routing (where the port choice determined for the first packet in each flow is used for all packets in the flow) according to the ordering requirements of the traffic.

(151) As has been described herein, a switch, such as switch **202** supports minimal and non-minimal path routing. It should be understood that minimal paths are based on the destination. If a destination NIC is local, an output port that connects to the destination switch is selected. If the destination is in another group, the packet is routed to a switch within the local group that is connected to the destination group. In a large system, there may only be one such path, but in a small system there are likely to be many, some connected to the input switch and others to switches elsewhere within its group. The input switch selects between them.

(152) The candidate ports considered for minimal path routing are further divided into preferred and nonpreferred sets of ports where ports in the preferred set may lead to a path containing fewer hops. Non-minimal paths route packets via an intermediary switch, referred to as a root switch. Root switches are selected on a packet-by-packet basis or a flow-by-flow basis depending on the ordering requirements of the traffic.

(153) Non-minimal traffic is routed “up” to the root switch, and then minimally “down” to the destination. In some embodiments, intermediate root switches are selected at random so as to distribute load uniformly. The network, e.g., network **100**, provides control over intermediate group selection enabling traffic to be routed towards intermediate groups that are known to be lightly loaded or away from those that have specific function or are known to be heavily loaded. Root switches may be distributed over all groups, where a non-minimal path may detect a root switch in the source group, the destination group, or any intermediate group. Global non-minimal routes take an indirect path through a root switch in an intermediate group. These paths require two global hops, one from the source group to an intermediate group, and one from the intermediate group to the destination group. Global non-minimal paths require up to three local hops, one in each group. The maximum path length is normally five switch-to-switch hops, whatever the system size. Although unlikely, it is possible to employ up to two additional local hops to work around faults within the network.

(154) Minimal routing is to be preferred as the paths are shorter, and hence the load on the network is lower. However, minimal routing alone will result in poor performance on some traffic patterns, for example when all nodes in one group communicate with nodes in one other group. Achieving good performance across a wide range of traffic patterns requires a mix of minimal and non-minimal routing.

(155) At each hop along a frame's path, the routing modes that may be used to advance the frame along its next hop are controlled by the configuration of the FRF routing algorithm table **408**. When a frame is received at a switch input port, the types of paths along with the frame may be forwarded is determined: local minimal, global minimal, local non-minimal, and global non-minimal. The set of output ports to which the frame may be forwarded is determined by the type of paths allowed at that point.

(156) The types of paths that are allowed to be taken depends on where the frame is at along its journey between its ingress and egress ports of the network. The path types are as follows. Local minimal paths select links that are directly connected to the frame's destination switch and may be used when the frame is in its destination group. Global minimal paths may be used when the frame is not in its destination group and select either global links that directly connect to the frame's destination group or local links that connect to a switch that has working global links that directly connect to the frame's destination group. Local non-minimal paths may be used when the frame is in its destination group, or when the frame is in an intermediate group. Local non-minimal paths

select local links connected to other switches in the group without regard for the frame's destination. When in the destination group, it must be possible to reach the frame's destination switch within one more hop after taking the local non-minimal hop. When in an intermediate group, it must be possible to reach a switch with a working global link that connects to the frame's destination group within one more hop after taking the local non-minimal hop. Local links, that connect to switches from which this is not possible, must not be selected.

(157) Global non-minimal-paths may be used when the frame is in its source group and is not in its destination group. Global non-minimal-paths select either global links connected to other groups or local links connected to other switches in the source group without regard for the frame's destination. Global links must only be selected if they connect to a group that has working links connecting it to the frame's destination group. Similarly, local links must only be selected if they connect to switches that have global links that are, themselves, valid global non-minimal path choices.

(158) Adaptive routing selects between minimal and non-minimal paths (described above) according to their current load.

(159) In terms of minimal routing, when in the destination group, but not at the destination switch, local minimal routes are generated by looking up the switch_id field of the destination fabric address in a local switch minimal table (FIG. 4A). The lookup returns a set of valid links. The local switch minimal table contains 128 entries each of 64 bits, with each bit representing one possible output port. When at the destination switch, the egress port or port choices are generated by looking up the endpoint_id field of the destination fabric address in the local port minimal table. The lookup returns a set of valid links. The local port minimal table contains 64 entries each of 64 bits, with each bit representing one possible output port.

(160) Global minimal routes are generated by looking up the global_id field of the destination fabric address in a global minimal table (FIG. 4A). The lookup returns a set of valid links. The global minimal table contains 512 entries each of 64 bits, with each bit representing one possible output port.

(161) Local minimal paths take at most one switch-to-switch hop, from the source switch to the destination switch, both of which are within the same group. There can be several such paths. Local non-minimal paths take two switch-to-switch hops, from the source switch to an intermediate switch, known as the Root switch, and from there to the destination switch. There are many such paths.

(162) Global minimal paths take one global hop from the source group to the destination group. There is at most, one local hop in each of the source and destination groups. Global minimal paths require a maximum of three switch-to-switch hops whatever the system size.

(163) In certain system configuration in which there are multiple global links connecting a source group to a destination group, a bias can arise such that the proportion of traffic injected at the source group that is distributed to each of the global links is not equal. As an example, consider the case of switch A, B, and C, all in the group X, with switch B having three global links connecting it to group V and switch C having two global links connecting it to group V. If traffic injected at switch A, destined for group V, is equally distributed between switch B and C, each of the two global links of switch C will be more heavily loaded than each of the three global links of switch B.

(164) To enable the bias to be counteracted, the Global Minimal Table can be divided into several blocks, each of which is capable of generating a valid set of global routing choices for any destination group. On a frame by frame basis, the block that is used to service the request is pseudo-randomly selected by the FRF. Within each Global Minimal Table block instance, only a subset of the possible candidate ports that can be used to reach the destination group are populated. The subset is chosen in such a way so as to counteract the bias. The populated subsets can be varied across the different block instances such that all possible candidate ports are able to be used.

(165) A global minimal path between an edge port in one group and an edge port in another group

can require one, two, or three hops across fabric links. One hop if the switch containing the ingress edge port in the source group has a global link directly connecting it to the switch containing the egress edge port in the destination group. Two hops if the frame traverses a global link between the two groups that is connected directly to either the ingress switch in the source group or to the egress switch in the destination group. In this case, one hop across a local link in either the source group or the destination group is also required. Lastly, three hops are required if the frame traverses a global link that is not directly connected to either the ingress switch in the source group or to the egress switch in the destination group. In this case, a hop across a local link is also required in both the source and destination groups. The local link hop in the source group takes the frame from the ingress switch to the source group switch that is connected to the global link. The local link hop in the destination group takes the frame from the destination group switch that is connected to the global link to the egress switch.

(166) When identifying minimal path candidate ports, the FRF is able to classify the candidates into a set of preferred ports and a set of non-preferred ports. The preferred ports are those that allow a global minimal path requiring two, or fewer fabric link hops. The non-preferred ports are all of the minimal path candidates that are not classified as preferred. Use of preferred paths, when available, and when not too heavily loaded reduces the average load on the system's local fabric links as it reduces the average number of local fabric links traversed per frame. Use of a preferred path may also reduce the end-to-end fabric latency experienced by the frame.

(167) When performing local non-minimal routing, any local link can be a candidate. However, some local links may need to be excluded from consideration if they lead to a switch from which it may not be possible to reach the destination because of link or switch failures that exist within the system. A CSR controls which ports are candidates for local non-minimal path routing.

(168) When performing global non-minimal routing, generally any global link can be a candidate. Additionally, generally any local link that reaches a switch with operational global links can also be a candidate. However, some links may need to be blocked from consideration if they lead to a switch or to a group from which it may not be possible to reach the destination group because of link or switch failures that exist within the system. CSRs control which ports are candidates for global nonminimal path routing.

(169) When selecting a candidate port to use for global non-minimal routing, if all candidate ports, global and local, are equally likely to be selected, for many system configurations globally non-minimal traffic will not be evenly distributed among the global links leaving the group. Consider, for example, the situation of three switches, A, B, and C within a group, where each switch is connected to each other switch by four local links, and switch A and switch B each terminate 14 global links and switch C terminates 16 global links. Ingress traffic arriving at an edge port of switch A may be routed globally non-minimally to any of the global or local links terminated by switch A. If it is equally likely to be routed to any of these links, then each global link terminated by switch B will receive only 4/14th (4 local links reaching 14 global links) of the traffic routed to each global link terminated by switch A. Similarly, each global link terminated by switch C will receive 4/16th (4 local links reaching 16 global links).

(170) To counteract this potential bias in the distribution of globally non-minimal traffic among a group's global links, when the set of global non-minimal candidate ports is being pseudo-randomly down selected to the small number of ports that will participate in the adaptive routing stage, a weighting can be applied to each of the candidate ports such that some will be more likely than others to survive the down-selection process.

(171) Adaptive routing selects between these minimal and non-minimal paths based on their load. Adaptively routed traffic starts on a minimal path, diverting to a non-minimal path if the load on the minimal path is high (this is known as progressive adaptive routing). Such paths are said to have diverged.

(172) A non-minimal path can be selected at the injection point or at the exit router in a source

group. A local non-minimal route may be taken within the source group, an intermediate group, or a destination group. Dragonfly routing algorithms allow a non-minimal path in both the intermediate and destination groups (consider a case where all traffic incoming on the global ports of a particular switch is destined for NICs on another switch in the group). In general however, non-minimal traffic is sufficiently well distributed as to avoid this happening, but an additional hop in the destination group may still be beneficial in cases where there is an error on a local link. Having arrived at the intermediate group the packet may take either a minimal route to the destination group or a local non-minimal route to a switch with a path to the destination group.

(173) Again, this decision is made based on load. Having taken a hop within the intermediate group the packet must detect root and take a minimal path to the destination. Adaptive decisions are made based on load and a bias towards preferred minimal, non-preferred minimal, or non-minimal. A routing algorithm, as described above, increases the bias towards minimal paths the closer a packet is to its destination. This algorithm prefers a direct path across the intermediate group provided the load is low.

(174) Restricted routing is used at points other than injection and root detection to prevent packets from flowing back in the direction from which they came. In a switch, cases in which a packet has taken one hop from points of injection and root detection are detected and ensured that a global port is taken. For local minimal routing, having taken a hop from the root, the packet will arrive at a switch that is connected to the destination NIC. For global minimal routing, having taken a hop from the point of injection, the packet will arrive at a switch with a global link, which must be taken. In the intermediate group, packets are allowed to take a local hop at the point of injection. From that point onward, the packet must follow a minimal path to its destination.

(175) When passing information from switch to switch, it is necessary to make instantaneous decisions about the next hop on the path. Decisions are taken using information derived from local state and information that is communicated from neighboring switches. Use of information from many different sources allows for more accurate/effective decisions. This includes information from neighbors.

(176) Prior systems carried information on average load from switch to switch. That said, more detailed information from related or neighboring switches is more helpful. In the current switch ASIC, a set of values can include information indicative of the status of output ports of related/neighboring switches. By passing this set of values, much better routing decisions can be made. In one example, a flag is passed back from neighboring switches, with the flag having one bit for each output port. For example, with a switch having 64 outputs, a 64 bit flag would be transmitted. This is much more accurate than simply passing a global average for neighboring ports.

(177) FIG. 7A illustrates an example scenario where average load is used as a basis for routing packets in switch fabric comprising source switch **702**, destination switch **708**, and two possible intermediate switches **704** and **706**. As illustrated in FIG. 7A, source switch **702**, based on adaptive routing as described herein, may determine two candidate intermediate switches (**704** and **706**) through which packets may be routed from source switch **702** to destination switch **708**. When basing adaptive routing decisions on average load of the intermediate switches **704** and **706**, in this example, intermediate switch **704** may have an average load value of two, while intermediate switch **706** may have an average load value of three. Basing adaptive routing decisions on average load value would then result in source switch **702** selecting to reach destination switch **708** by way of intermediate switch **704** (i.e., route or path **710**). However, the link between intermediate switch **704** and destination switch **708** may be busy, which as noted above may negatively impact, e.g., latency.

(178) In accordance with one embodiment, and as described above, port busyness can be taken into account. That is, and referring back to FIG. 5, down-selection of candidate ports and adaptive route selection can be based on load. Again, FRF component **400** considers three categories of candidate

ports to which a frame may be forwarded: preferred minimal path candidate ports **502**; non-preferred minimal path candidate ports **504**; and non-minimal path candidate ports **506**. Depending on where a frame may be along its path, non-minimal path candidate ports are either global, non-minimal path candidate ports or local non-minimal path candidate ports.

(179) Filtering may be applied to the three categories of candidate ports, e.g., operational port filtering, useable port filtering, and busy port filtering. Port filtering as applied herein can be used to reduce the set of valid ports considered as path candidate ports by identifying and removing absent and/or faulty ports and/or busy ports from consideration.

(180) In the example of FIG. 7B, busy port filtering may result in a different path being selected to route frames from source switch **702** to destination switch **708**. Again, knowledge regarding which ports are busy in a neighboring switch (in this example, intermediate switches **704** and **706**) can be used to determine if the ports that connect to a neighboring switch are poor candidates to receive a frame based on whether the neighboring switch will subsequently need to forward the frame to a port that is already busy. Accordingly, busy port filtering can be performed by FRF component **400** by using busy port masks to remove heavily loaded ports from being considered candidate ports. FRF component **400** maintains four masks of busy ports, that is, ports whose load exceeds a software-defined threshold: local switch busy port mask; global non-minimal busy global port mask; global non-minimal busy local port mask; remote switch busy port mask. Information from these masks is communicated between switches to populate the remote switch.

(181) Here, a local switch busy port mask can be applied to minimal path candidate ports as well as to local non-minimal path candidate ports, e.g., output ports of intermediate switches **704** and **706**. FRF component **400** generates a 64-bit `Is_busy_port_mask` by comparing each port's `local_load` to a software defined threshold. Ports with loads higher than this threshold are marked as busy in this mask. In this example, output ports from intermediate switch **704** directed to destination switch **708** meet or surpass the software defined threshold specified for the switch fabric. Therefore, the `Is_busy_port_mask` indicates that the link from intermediate switch **704** to destination switch **708** is busy. On the other hand, the output ports of intermediate switch **706** to destination switch **708** are not marked or flagged with the `Is_busy_port_mask` as the busyness threshold is not met/surpassed, reflecting that the link between intermediate switch **706** and destination switch **708** is quiet, and not busy.

(182) Based on the busy port mask indication that the link between intermediate switch **704** and **708** is busy, the output port of intermediate switch **704** linked to an input port of destination switch **708** is deemed to be a poor candidate, and removed from consideration. Accordingly, and in contrast to the scenario of FIG. 7A, where route **710** is selected based on average load, route or path **712** is selected to route a frame(s) from source switch **702** through intermediate switch **706**, and on to its destination, i.e., destination switch **708**. It should be understood that the quiet link condition between intermediate switch **706** and destination switch **708** is a better choice to route frames from source **702** despite the average load being higher.

(183) Similarly, as noted above, FRF component **400** can use received remote switch busy port masks to generate the aforementioned Remote Switch Busy Global Port Table and Remote Switch Busy Local Port Table of busy port masks that identify ports connected to neighboring switches that should be avoided as global minimal path candidates based on the destination group to which the frame is being routed and that should be avoided as local non-minimal path candidates based on the destination switch to which the frame is being routed. Thus, the `rs_busy_port_mask` can also be used to identify ports connected to neighboring switches that may be busy or quiet in accordance with some embodiments. In the example, of FIG. 7B, an input port of destination switch **708** may be deemed busy by way of `rs_busy_port_masks` applied thereto. Such ports may be avoided as local non-minimal path candidates, when routing in the destination switch, based on the destination switch to which the frame is being routed

(184) FIG. 8 is a flow chart illustrating example operations that may be performed to effectuate

adaptive route or path selection in accordance with one embodiment of the disclosed technology. As noted above, switch **202** may be an ASIC, and thus one or more electronic circuits that include electronic components are provided for controlling or instantiating FRF component **400** (FIGS. **4A** and **4B**) by performing the disclosed example operations.

(185) At operation **800**, a plurality of load signals from a plurality of neighboring switches is received, each of the plurality of load signals comprising a set of values indicative of a load at each of the plurality of neighboring switches providing the load signal, each value within the set of values providing an indication for each link of the plurality of links attached thereto as to whether the link is busy or quiet. As described above, a switch, e.g., switch **202** may have a plurality of ports (**64** ports in some embodiments). Packets may be routed between input and output ports of one switch to another one or more switches (each having their own input and output ports. Routing of packets between switches can be performed in accordance with certain characteristics of the switches that the packets may traverse along a path from a source switch to a destination switch.

(186) Between the source and destination switches may be one or more intermediate switches. In order to select a path, ports of switches over which the packet may traverse are filtered to identify certain path candidate ports that have desirable characteristics regarding, e.g., how busy they are in terms of traffic traversing through the ports. Accordingly, adaptive routing is effectuated by determining, based on load (an indication of busyness) associated with a candidate port, whether or not the packet should traverse path including a candidate port of a switch, such as an intermediate switch. In some embodiments, the local load is compared to a threshold that can be specified by the CSR configuration, and ports can be excluded based on whether or not the local load falls below or surpasses that threshold.

(187) At operation **802**, based upon the plurality of load signals, an output link is selected for routing a received packet and routing the received packet via the selected output link. That is, a mask can be applied to those ports whose local load surpasses the threshold, i.e., those ports are marked as being busy in this mask. It should be understood that other busy port masks can be applied in accordance with other embodiments, e.g., global non-minimal busy global port mask, global non-minimal busy local port mask, or a destination group dependent busy-port mask depending on the circumstances. For example, a destination group dependent busy-port mask, obtained from a remote switch busy global port table can be applied to global minimal path candidate ports. Correspondingly, when a frame is being routed in its destination group, a destination switch-dependent busy-port mask, obtained from a remote switch busy local port table can be applied to local non-minimal path candidate ports. Thus, the use of load information need not be limited to simply determining ports of intermediate switches. Upon a port being marked as busy in a busy port mask, that port may be removed or taken out of consideration as a candidate port. As illustrated in FIG. **5**, and described above, load-based adaptive selection can narrow surviving candidate ports to a least loaded candidate port. A path along which the packet is routed can be selected to include this least loaded port, and the packet is ultimately routed through the selected one of the plurality of switch ports.

(188) It should be noted that the terms “optimize,” “optimal” and the like as used herein can be used to mean making or achieving performance as effective or perfect as possible. However, as one of ordinary skill in the art reading this document will recognize, perfection cannot always be achieved. Accordingly, these terms can also encompass making or achieving performance as good or effective as possible or practical under the given circumstances, or making or achieving performance better than that which can be achieved with other settings or parameters.

(189) FIG. **9** depicts a block diagram of an example computer system **900** in which various of the embodiments described herein may be implemented. The computer system **900** includes a bus **902** or other communication mechanism for communicating information, one or more hardware processors **904** coupled with bus **902** for processing information.

(190) Hardware processor(s) **904** may be, for example, one or more general purpose

microprocessors.

(191) The computer system **900** also includes a main memory **906**, such as a random access memory (RAM), cache and/or other dynamic storage devices, coupled to bus **902** for storing information and instructions to be executed by processor **904**. Main memory **906** also may be used for storing temporary variables or other intermediate information during execution of instructions to be executed by processor **904**. Such instructions, when stored in storage media accessible to processor **904**, render computer system **900** into a special-purpose machine that is customized to perform the operations specified in the instructions.

(192) The computer system **900** further includes a read only memory (ROM) **908** or other static storage device coupled to bus **902** for storing static information and instructions for processor **904**. A storage device **910**, such as a magnetic disk, optical disk, or USB thumb drive (Flash drive), etc., is provided and coupled to bus **902** for storing information and instructions.

(193) The computer system **900** may be coupled via bus **902** to a display **912**, such as a liquid crystal display (LCD) (or touch screen), for displaying information to a computer user. An input device **914**, including alphanumeric and other keys, is coupled to bus **902** for communicating information and command selections to processor **904**. Another type of user input device is cursor control **916**, such as a mouse, a trackball, or cursor direction keys for communicating direction information and command selections to processor **904** and for controlling cursor movement on display **912**. In some embodiments, the same direction information and command selections as cursor control may be implemented via receiving touches on a touch screen without a cursor.

(194) The computing system **900** may include a user interface module to implement a GUI that may be stored in a mass storage device as executable software codes that are executed by the computing device(s). This and other modules may include, by way of example, components, such as software components, object-oriented software components, class components and task components, processes, functions, attributes, procedures, subroutines, segments of program code, drivers, firmware, microcode, circuitry, data, databases, data structures, tables, arrays, and variables.

(195) In general, the word “component,” “engine,” “system,” “database,” “data store,” and the like, as used herein, can refer to logic embodied in hardware or firmware, or to a collection of software instructions, possibly having entry and exit points, written in a programming language, such as, for example, Java, C or C++. A software component may be compiled and linked into an executable program, installed in a dynamic link library, or may be written in an interpreted programming language such as, for example, BASIC, Perl, or Python. It will be appreciated that software components may be callable from other components or from themselves, and/or may be invoked in response to detected events or interrupts. Software components configured for execution on computing devices may be provided on a computer readable medium, such as a compact disc, digital video disc, flash drive, magnetic disc, or any other tangible medium, or as a digital download (and may be originally stored in a compressed or installable format that requires installation, decompression or decryption prior to execution). Such software code may be stored, partially or fully, on a memory device of the executing computing device, for execution by the computing device. Software instructions may be embedded in firmware, such as an EPROM. It will be further appreciated that hardware components may be comprised of connected logic units, such as gates and flip-flops, and/or may be comprised of programmable units, such as programmable gate arrays or processors.

(196) The computer system **900** may implement the techniques described herein using customized hard-wired logic, one or more ASICs or FPGAs, firmware and/or program logic which in combination with the computer system causes or programs computer system **900** to be a special-purpose machine. According to one embodiment, the techniques herein are performed by computer system **900** in response to processor(s) **904** executing one or more sequences of one or more instructions contained in main memory **906**. Such instructions may be read into main memory **906**

from another storage medium, such as storage device **910**. Execution of the sequences of instructions contained in main memory **906** causes processor(s) **904** to perform the process steps described herein. In alternative embodiments, hard-wired circuitry may be used in place of or in combination with software instructions.

(197) The term “non-transitory media,” and similar terms, as used herein refers to any media that store data and/or instructions that cause a machine to operate in a specific fashion. Such non-transitory media may comprise non-volatile media and/or volatile media. Non-volatile media includes, for example, optical or magnetic disks, such as storage device **910**. Volatile media includes dynamic memory, such as main memory **906**. Common forms of non-transitory media include, for example, a floppy disk, a flexible disk, hard disk, solid state drive, magnetic tape, or any other magnetic data storage medium, a CD-ROM, any other optical data storage medium, any physical medium with patterns of holes, a RAM, a PROM, and EPROM, a FLASH-EPROM, NVRAM, any other memory chip or cartridge, and networked versions of the same.

(198) Non-transitory media is distinct from but may be used in conjunction with transmission media. Transmission media participates in transferring information between non-transitory media. For example, transmission media includes coaxial cables, copper wire and fiber optics, including the wires that comprise bus **902**. Transmission media can also take the form of acoustic or light waves, such as those generated during radio-wave and infra-red data communications.

(199) The computer system **900** also includes a communication interface **918** coupled to bus **902**. Network interface **918** provides a two-way data communication coupling to one or more network links that are connected to one or more local networks. For example, communication interface **918** may be an integrated services digital network (ISDN) card, cable modem, satellite modem, or a modem to provide a data communication connection to a corresponding type of telephone line. As another example, network interface **918** may be a local area network (LAN) card to provide a data communication connection to a compatible LAN (or WAN component to communicated with a WAN). Wireless links may also be implemented. In any such implementation, network interface **918** sends and receives electrical, electromagnetic or optical signals that carry digital data streams representing various types of information.

(200) A network link typically provides data communication through one or more networks to other data devices. For example, a network link may provide a connection through local network to a host computer or to data equipment operated by an Internet Service Provider (ISP). The ISP in turn provides data communication services through the world wide packet data communication network now commonly referred to as the “Internet.” Local network and Internet both use electrical, electromagnetic or optical signals that carry digital data streams. The signals through the various networks and the signals on network link and through communication interface **918**, which carry the digital data to and from computer system **900**, are example forms of transmission media.

(201) The computer system **900** can send messages and receive data, including program code, through the network(s), network link and communication interface **918**. In the Internet example, a server might transmit a requested code for an application program through the Internet, the ISP, the local network and the communication interface **918**.

(202) The received code may be executed by processor **904** as it is received, and/or stored in storage device **910**, or other non-volatile storage for later execution.

(203) Each of the processes, methods, and algorithms described in the preceding sections may be embodied in, and fully or partially automated by, code components executed by one or more computer systems or computer processors comprising computer hardware. The one or more computer systems or computer processors may also operate to support performance of the relevant operations in a “cloud computing” environment or as a “software as a service” (SaaS). The processes and algorithms may be implemented partially or wholly in application-specific circuitry. The various features and processes described above may be used independently of one another, or may be combined in various ways. Different combinations and sub-combinations are intended to

fall within the scope of this disclosure, and certain method or process blocks may be omitted in some implementations. The methods and processes described herein are also not limited to any particular sequence, and the blocks or states relating thereto can be performed in other sequences that are appropriate, or may be performed in parallel, or in some other manner. Blocks or states may be added to or removed from the disclosed example embodiments. The performance of certain of the operations or processes may be distributed among computer systems or computers processors, not only residing within a single machine, but deployed across a number of machines. (204) As used herein, a circuit might be implemented utilizing any form of hardware, software, or a combination thereof. For example, one or more processors, controllers, ASICs, PLAs, PALs, CPLDs, FPGAs, logical components, software routines or other mechanisms might be implemented to make up a circuit. In implementation, the various circuits described herein might be implemented as discrete circuits or the functions and features described can be shared in part or in total among one or more circuits. Even though various features or elements of functionality may be individually described or claimed as separate circuits, these features and functionality can be shared among one or more common circuits, and such description shall not require or imply that separate circuits are required to implement such features or functionality. Where a circuit is implemented in whole or in part using software, such software can be implemented to operate with a computing or processing system capable of carrying out the functionality described with respect thereto, such as computer system **900**.

(205) As used herein, the term “or” may be construed in either an inclusive or exclusive sense. Moreover, the description of resources, operations, or structures in the singular shall not be read to exclude the plural. Conditional language, such as, among others, “can,” “could,” “might,” or “may,” unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps.

(206) Terms and phrases used in this document, and variations thereof, unless otherwise expressly stated, should be construed as open ended as opposed to limiting. Adjectives such as “conventional,” “traditional,” “normal,” “standard,” “known,” and terms of similar meaning should not be construed as limiting the item described to a given time period or to an item available as of a given time, but instead should be read to encompass conventional, traditional, normal, or standard technologies that may be available or known now or at any time in the future. The presence of broadening words and phrases such as “one or more,” “at least,” “but not limited to” or other like phrases in some instances shall not be read to mean that the narrower case is intended or required in instances where such broadening phrases may be absent.

Claims

1. A method of passing data amongst a plurality of switches, the method comprising: at a switch, receiving a set of values indicating, for each link of a plurality of links attached to a neighboring switch as to whether the link is busy or quiet, the switch and the neighboring switch being part of a switch fabric, wherein the indication as to whether the link is busy comprises a remote switch busy port mask characterizing ports of the neighboring switches, one or more of the remote switch busy port masks being transformed to form entries of a remote switch busy global port table and one or more of the remote switch busy port masks being transformed to form entries of a remote switch busy local port table, the remote switch busy port masks being received from the neighboring switches; and based upon the set of values, selecting an output link for routing a received packet and routing the received packet through the switch fabric via the selected output link.
2. The method of claim 1, wherein the switch comprises a source switch, and wherein the neighboring switch comprises one of a destination switch and an intermediate switch.
3. The method of claim 2, further comprising selecting a path over which the data is routed.

4. The method of claim 3, further comprising filtering ports of the source switch and at least one of the intermediate switch and destination switch, and identifying path candidate ports having desirable traffic characteristics.
5. The method of claim 4, wherein the set of values indicating whether the link is busy or quiet is compared to a configuration-specific threshold.
6. The method of claim 5, further comprising, applying a busy mask to those ports making up a link that is busy.
7. The method of claim 4, wherein the filtering of the ports comprises filtering into one of three candidate port types including operational ports, useable ports, and busy ports.
8. A switch, comprising: an application-specific integrated circuit (ASIC) to: receive an indication, for each link of a plurality of links attached to a neighboring switch as to whether the link is busy or quiet, the switch and the neighboring switch comprising a switch fabric; applying a busy port mask to those input and output ports making up each of the links identified as being busy, and taking those input and output ports out of consideration as candidate ports; of those input and output ports making up each of the links identified as being quiet, selecting an output link comprised of quiet input and output ports for routing a received packet and routing the received packet through the switch fabric via the selected output link.
9. The switch of claim 8, wherein the switch comprises a source switch, and wherein the neighboring switch comprises one of a destination switch and an intermediate switch.
10. The switch of claim 9, further comprising selecting a path over which data is routed.
11. The switch of claim 10, further comprising filtering ports of the source switch and at least one of the intermediate switch and destination switch, and identifying path candidate ports having desirable traffic characteristics.
12. The switch of claim 11, wherein the indication whether the link is busy or quiet is compared to a configuration-specific threshold.
13. The switch of claim 11, wherein the filtering of the ports comprises filtering into one of three candidate port types including operational ports, useable ports, and busy ports.
14. A system, comprising: a plurality of switches interconnected to form a switch fabric; at a first switch of the plurality of switches, receiving an indication, for each link of a plurality of links attached to a plurality of neighboring switches of the plurality of switches as to whether the link is busy or quiet, wherein the indication as to whether the link is busy comprises a remote switch busy port mask characterizing ports of the neighboring switches, one or more of the remote switch busy port masks being transformed to form entries of a remote switch busy global port table and one or more of the remote switch busy port masks being transformed to form entries of a remote switch busy local port table, the remote switch busy port masks being received from the neighboring switches; and based upon the indication, selecting an output link for routing a received packet and routing the received packet through the switch fabric via the selected output link.
15. The system of claim 14, wherein the remote switch busy global port table is indexed in accordance with destination group identifiers corresponding to those of the destination groups of the received packet.
16. The system of claim 15, wherein subsets of the neighboring switches comprise global minimal paths to reach the destination groups.
17. The system of claim 15, wherein a fabric routing function instance of the first switch, in accordance with an executed routing algorithm, determines a recommended port to which the received packet is to be routed subsequent to filtering out candidate ports based on a determination that the filtered out candidate ports are busy in accordance with the indexed entry in the remote switch busy global port table.
18. The system of claim 14, wherein the remote switch busy local port table is indexed in accordance with switch identifiers corresponding to those of the destination switches of the received packets.

19. The system of claim 18, wherein subsets of the neighboring switches comprise local non-minimal and local minimal paths to the destination switches.

20. The system of claim 18, wherein a fabric routing function instance of the first switch, in accordance with an executed routing algorithm, determines a recommended port to which the received packet is to be routed subsequent to filtering out candidate ports based on a determination that the filtered out candidate ports are busy in accordance with the indexed entry in the remote switch busy local port table.
