

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386774
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Balachandran; Jayaprakash et al.

AI accelerator apparatus using full mesh connectivity chiplet devices for transformer workloads

Abstract

An AI accelerator apparatus using in-memory compute chiplet devices. The apparatus includes a first semiconductor substrate having a plurality of chiplets, each of which includes a plurality of tiles. Each tile includes a plurality of slices, a central processing unit (CPU), and a hardware dispatch device. Each slice can include a digital in-memory compute (DIMC) device configured to perform high throughput computations. In particular, the DIMC device can be configured to accelerate the computations of attention functions for transformer-based models (a.k.a. transformers) applied to machine learning applications. The chiplets are in a full mesh connectivity configuration such that at least one of the die-to-die (D2D) interconnects of each chiplet is coupled to one of the D2D interconnects of each other chiplet using a non-diagonal link. The chiplets can also include other interfaces to facilitate communication between the chiplets, memory and a server or host system.

Inventors: Balachandran; Jayaprakash (Santa Clara, CA), Ankit; Aayush (Santa Clara, CA), Quek; Irene (San Jose, CA), Bhoja; Sudeep (Cupertino, CA)

Applicant: d-MATRIX CORPORATION (Santa Clara, CA)

Family ID: 95340553

Assignee: d-MATRIX CORPORATION (Santa Clara, CA)

Appl. No.: 18/486872

Filed: October 13, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250123984 A1	Apr. 17, 2025

Publication Classification

Int. Cl.: G06F13/42 (20060101); G06F9/38 (20180101)

U.S. Cl.:

CPC G06F13/4221 (20130101); G06F9/38873 (20230801);

Field of Classification Search

CPC: G06F (13/4221); G06F (9/38873)

USPC: 710/313

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11848670	12/2022	Serra	N/A	H03K 19/1776
11977766	12/2023	Dally	N/A	G06F 3/0655
12013808	12/2023	Koker	N/A	G06F 17/16
12099453	12/2023	Dally	N/A	G06F 13/1689
12260118	12/2024	Maheshwari	N/A	G11C 11/4096
2019/0245582	12/2018	Burton	N/A	N/A
2020/0082246	12/2019	Shao	N/A	G06N 3/048
2020/0151847	12/2019	Schluessler	N/A	G06T 17/10
2022/0100247	12/2021	Garg	N/A	G06F 1/3203
2022/0107808	12/2021	Enamandram	N/A	G06F 9/3863
2022/0222194	12/2021	Chandwani	N/A	G06F 13/1663
2022/0327083	12/2021	Das Sharma et al.	N/A	N/A
2022/0334995	12/2021	Das Sharma et al.	N/A	N/A
2023/0168899	12/2022	Bhoja et al.	N/A	N/A
2023/0169021	12/2022	Bhoja et al.	N/A	N/A
2024/0143505	12/2023	Misra	N/A	G06F 12/0811
2024/0168915	12/2023	Du	N/A	G06F 15/825
2025/0029659	12/2024	Karda	N/A	G11C 16/08

OTHER PUBLICATIONS

International Search Report and Written Opinion for related application PCT/US2024/050581 mailed Jan. 17, 2025. cited by applicant

Primary Examiner: Shin; Christopher B

Attorney, Agent or Firm: Ogawa P.C.

Background/Summary

CROSS-REFERENCES TO RELATED APPLICATIONS

(1) N/A

BACKGROUND OF THE INVENTION

(2) The present invention relates generally to integrated circuit (IC) devices and artificial intelligence (AI). More specifically, the present invention relates to methods and device structures for accelerating computing workloads in transformer-based models (a.k.a. transformers).

(3) The transformer has been the dominant neural network architecture in the natural language processing (NLP) field, and its use continues to expand into other machine learning applications. The original Transformer was introduced in the paper “Attention is all you need” (Vaswani et al., 2017), which sparked the development of many transformer model variations, such as the generative pre-trained transformer (GPT) and the bidirectional encoder representations from transformers (BERT) models. Such transformers have significantly outperformed other models in inference tasks by their use of a self-attention mechanism that avoids recursion and allows for easy parallelism. On the other hand, the transformer workloads are very computationally intensive and have high memory requirements, and have been plagued as being time-intensive and inefficient.

(4) Most recently, NLP models have grown by a thousand times in both model size and compute requirements. For example, it can take about 4 months for 1024 graphics processing units (GPUs) to train a model like GPT-3 with 175 billion parameters. New NLP models having a trillion parameters are already being developed, and multi-trillion parameter models are on the horizon. Such rapid growth has made it increasingly difficult to serve NLP models at scale.

(5) From the above, it can be seen that improved devices and method to accelerate compute workloads for transformers are highly desirable.

BRIEF SUMMARY OF THE INVENTION

(6) The present invention relates generally to integrated circuit (IC) devices and artificial intelligence (AI) systems. More particularly, the present invention relates to methods and device structures for accelerating computing workloads in transformer-based neural network models (a.k.a. transformers). These methods and structures can be used in machine/deep learning applications such as natural language processing (NLP), computer vision (CV), and the like. Merely by way of example, the invention has been applied to AI accelerator apparatuses and chiplet devices configured to perform high throughput operations for NLP.

(7) According to an example, the present invention provides for a method and structure of an AI accelerator apparatus configured with in-memory compute and full mesh connectivity. The apparatus can include a plurality of chiplets coupled together in the full mesh connectivity configuration. Each of these chiplets can include a plurality of tiles, and each of these tiles can include at least a plurality of slices, a central processing unit (CPU) coupled to the plurality of slices, and a hardware dispatch device coupled to the CPU. The apparatus includes a plurality of die-to-die (D2D) interconnects coupled to each of the CPUs in each of the tiles, and at least one of the D2D interconnects of each chiplet is coupled to one of the D2D interconnects of each other chiplet using a non-diagonal link. In a specific example, the plurality of chiplets includes at least four chiplets, and each of the chiplets is coupled to each other chiplet in the full mesh connectivity configuration using a plurality of intra-chiplet non-diagonal links and a plurality of inter-chiplet non-diagonal links.

(8) The apparatus can also include a first clock configured to output a clock signal of about 0.5 GHz to 4 GHz, and each of the slices can include a digital in memory compute (DIMC) device coupled to a second clock configured at an output rate of one half of the rate of the first clock. As discussed previously, this DIMC device can be configured to allow for a throughput of one or more matrix computations provided in the DIMC device such that the throughput is characterized by 512 multiply accumulates per a clock cycle. The chiplets can also include interconnect interfaces (e.g., PCIe interfaces, or the like), memory interfaces (e.g., DRAM interfaces, or the like), global CPU

interfaces (e.g., RISC interfaces, or the like), as well as other interfaces to facilitate communication between the chiplets, memory and a server or host system.

(9) The AI accelerator and chiplet device architecture and its related methods can provide many benefits. With modular chiplets, the AI accelerator apparatus can be easily scaled to accelerate the workloads for transformers of different sizes. The DIMC configuration within the chiplet slices also improves computational performance and reduces power consumption by integrating computational functions and memory fabric. Further, embodiments of the AI accelerator apparatus can allow for quick and efficient mapping from the transformer to enable effective implementation of AI applications.

(10) A further understanding of the nature and advantages of the invention may be realized by reference to the latter portions of the specification and attached drawings.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) In order to more fully understand the present invention, reference is made to the accompanying drawings. Understanding that these drawings are not to be considered limitations in the scope of the invention, the presently described embodiments and the presently understood best mode of the invention are described with additional detail through use of the accompanying drawings in which:

(2) FIG. 1A-1B are simplified block diagrams illustrating AI accelerator apparatuses according to examples of the present invention.

(3) FIGS. 2A-2D are simplified block diagrams illustrating 16-slice chiplet devices according to examples of the present invention.

(4) FIGS. 3A-B are simplified block diagrams illustrating slice devices according to examples of the present invention.

(5) FIG. 4 is a simplified block diagram illustrating an in-memory-compute (IMC) module according to an example of the present invention.

(6) FIG. 5A is a simplified block flow diagram illustrating numerical formats of the data being processed in a slice device according to an example of the present invention.

(7) FIG. 5B is a simplified diagram illustrating example numerical formats.

(8) FIG. 6 is a simplified block diagram of a transformer architecture.

(9) FIG. 7 is a simplified diagram illustrating a self-attention layer process for an example NLP model.

(10) FIG. 8 is a simplified block diagram illustrating an example transformer.

(11) FIG. 9 is a simplified block diagram illustrating an attention head layer of an example transformer.

(12) FIG. 10 is a simplified table representing an example mapping process between a 24-layer transformer and an example eight-chiplet AI accelerator apparatus according to an example of the present invention.

(13) FIG. 11 is a simplified block flow diagram illustrating a mapping process between a transformer and an AI accelerator apparatus according to an example of the present invention.

(14) FIG. 12 is a simplified table representing a tiling attention process of a transformer to an AI accelerator apparatus according to an example of the present invention.

(15) FIGS. 13A-13C are simplified tables illustrating data flow through the IMC and single input multiple data (SIMD) modules according to an example of the present invention.

(16) FIGS. 14A and 14B shows a simplified diagram illustrating a method of forming a stacked wafer apparatus according to an example of the present invention.

(17) FIG. 14C is a simplified block diagram illustrating a memory device according to an example of the present invention.

- (18) FIG. 15 is a simplified block illustrating a 3D stacked device die according to an example of the present invention.
- (19) FIGS. 16A and 16B are simplified diagrams illustrating an error correction implementation for a memory device according to an example of the present invention.
- (20) FIGS. 17A-17C are simplified diagrams illustrating a memory tiling configuration for a stacked wafer apparatus according to an example of the present invention.
- (21) FIG. 18 is a simplified diagram illustrating a cross-sectional view of a stacked logic and memory device according to an example of the present invention.
- (22) FIGS. 19A-19C are simplified diagrams illustrating a 3D stacked chiplet and memory device using a hybrid-bonding interconnection according to various examples of the present invention.
- (23) FIG. 19D is a device level image showing a cross-sectional view of a 3D stacked chiplet and memory device according to an example of the present invention.
- (24) FIG. 20 is a simplified diagram illustrating a top view of a 3D stacked AI accelerator apparatus according to an example of the present invention.
- (25) FIGS. 21A-21F are simplified diagrams illustrating cross-sectional views of 3D stacked chiplet devices according to various examples of the present invention.
- (26) FIG. 22 is a simplified diagram illustrating a top view of an example AI accelerator apparatus using full mesh chiplet connectivity according to an example of the present invention.
- (27) FIG. 23 is a simplified flow diagram illustrating a method of operating a transformer model according to an example of the present invention.
- (28) FIG. 24 is a simplified graph illustrating the effect of memory bandwidth and interconnect bandwidth on performance of processing transformer workloads.
- (29) FIG. 25 is a simplified graph illustrating the evolution of byte/flop (B/F) ratio for deep neural networks.
- (30) FIG. 26 is a simplified graph illustrating the power efficiency of high bandwidth memory (HBM).
- (31) FIG. 27A is a simplified graph showing throughput data for 3D stacked AI accelerator apparatuses according to various examples of the present invention.
- (32) FIG. 27B is a simplified graph showing latency data for 3D stacked AI accelerator apparatuses according to various examples of the present invention.
- (33) FIG. 27C is a simplified graph showing a comparison of throughput per card, latency per token, and number of cards in the system for an HBM device configuration and two 3D stacked apparatus configurations according to examples of the present invention.
- (34) FIG. 28A is a simplified block diagram illustrating a 3D stacked AI accelerator apparatus with a memory die having memory bank group devices according to an example of the present invention.
- (35) FIG. 28B is a simplified block diagram illustrating a 3D stacked AI accelerator apparatus with multiple stacked memory dies having memory bank group devices according to an example of the present invention.
- (36) FIG. 28C is a simplified block diagram illustrating a memory bank group device configuration according to an example of the present invention.
- (37) FIG. 29A is a simplified diagram illustrating a cross-sectional view of a 3D stacked AI accelerator apparatus with an organic substrate according to an example of the present invention.
- (38) FIG. 29B is a simplified diagram illustrating a cross-sectional view of a 3D stacked AI accelerator apparatus with an organic substrate and an embedded capacitor die according to an example of the present invention.
- (39) FIG. 29C is a simplified diagram illustrating a cross-sectional view of a 3D stacked AI accelerator apparatus with an organic substrate and multiple stacked memory dies according to an example of the present invention.
- DETAILED DESCRIPTION OF THE INVENTION

(40) The present invention relates generally to integrated circuit (IC) devices and artificial intelligence (AI) systems. More particularly, the present invention relates to methods and device structures for accelerating computing workloads in transformer-based neural network models (a.k.a. transformers). These methods and structures can be used in machine/deep learning applications such as natural language processing (NLP), computer vision (CV), and the like. Merely by way of example, the invention has been applied to AI accelerator apparatuses and chiplet devices configured to perform high throughput operations for NLP.

(41) Currently, the vast majority of NLP models are based on the transformer model, such as the bidirectional encoder representations from transformers (BERT) model, BERT Large model, and generative pre-trained transformer (GPT) models such as GPT-2 and GPT-3, etc. However, these transformers have very high compute and memory requirements. According to an example, the present invention provides for an apparatus using chiplet devices that are configured to accelerate transformer computations for AI applications. Examples of the AI accelerator apparatus are shown in FIGS. 1A and 1B.

(42) FIG. 1A illustrates a simplified AI accelerator apparatus **101** with two chiplet devices **110**. As shown, the chiplet devices **110** are coupled to each other by one or more die-to-die (D2D) interconnects **120**. Also, each chiplet device **110** is coupled to a memory interface **130** (e.g., static random access memory (SRAM), dynamic random access memory (DRAM), synchronous dynamic RAM (SDRAM), or the like). The apparatus **101** also includes a substrate member **140** that provides mechanical support to the chiplet devices **110** that are configured upon a surface region of the substrate member **140**. The substrate can include interposers, such as a silicon interposer, glass interposer, organic interposer, or the like. The chiplets can be coupled to one or more interposers, which can be configured to enable communication between the chiplets and other components (e.g., serving as a bridge or conduit that allows electrical signals to pass between internal and external elements).

(43) FIG. 1B illustrates a simplified AI accelerator apparatus **102** with eight chiplet devices **110** configured in two groups of four chiplets on the substrate member **140**. Here, each chiplet device **110** within a group is coupled to other chiplet devices by one or more D2D interconnects **120**. Apparatus **102** also shows a DRAM memory interface **130** coupled to each of the chiplet devices **110**. The DRAM memory interface **130** can be coupled to one or more memory modules, represented by the “Mem” block.

(44) As shown, the AI accelerator apparatuses **101** and **102** are embodied in peripheral component interconnect express (PCIe) card form factors, but the AI accelerator apparatus can be configured in other form factors as well. These PCIe card form factors can be configured in a variety of dimensions (e.g., full height, full length (FHFL); half height, half length (HHHL), etc.) and mechanical sizes (e.g., 1×, 2×, 4×, 16×, etc.). In an example, one or more substrate members **140**, each having one or more chiplets, are coupled to a PCIe card. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to these elements and configurations of the AI accelerator apparatus.

(45) Embodiments of the AI accelerator apparatus can implement several techniques to improve performance (e.g., computational efficiency) in various AI applications. The AI accelerator apparatus can include digital in-memory-compute (DIMC) to integrate computational functions and memory fabric. Algorithms for the mapper, numerics, and sparsity can be optimized within the compute fabric. And, use of chiplets and interconnects configured on organic interposers can provide modularity and scalability.

(46) According to an example, the present invention implements chiplets with in-memory-compute (IMC) functionality, which can be used to accelerate the computations required by the workloads of transformers. The computations for training these models can include performing a scaled dot-product attention function to determine a probability distribution associated with a desired result in a particular AI application. In the case of training NLP models, the desired result can include

predicting subsequent words, determining contextual word meaning, translating to another language, etc.

(47) The chiplet architecture can include a plurality of slice devices (or slices) controlled by a central processing unit (CPU) to perform the transformer computations in parallel. Each slice is a modular IC device that can process a portion of these computations. The plurality of slices can be divided into tiles/gangs (i.e., subsets) of one or more slices with a CPU coupled to each of the slices within the tile. This tile CPU can be configured to perform transformer computations in parallel via each of the slices within the tile. A global CPU can be coupled to each of these tile CPUs and be configured to perform transformer computations in parallel via all of the slices in one or more chiplets using the tile CPUs. Further details of the chiplets are discussed in reference to FIGS. 2A-5B, while transformers are discussed in reference to FIGS. 6-9.

(48) FIG. 2A is a simplified block diagram illustrating an example configuration of a 16-slice chiplet device **201**. In this case, the chiplet **201** includes four tile devices **210**, each of which includes four slice devices **220**, a CPU **221**, and a hardware dispatch (HW DS) device **222**. In a specific example, these tiles **210** are arranged in a symmetrical manner. As discussed previously, the CPU **221** of a tile **210** can coordinate the operations performed by all slices within the tile. The HW DS **222** is coupled to the CPU **221** and can be configured to coordinate control of the slices **220** in the tile **210** (e.g., to determine which slice in the tile processes a target portion of transformer computations). In a specific example, the CPU **221** can be a reduced instruction set computer (RISC) CPU, or the like. Further, the CPU **221** can be coupled to a dispatch engine, which is configured to coordinate control of the CPU **221** (e.g., to determine which portions of transformer computations are processed by the particular CPU).

(49) The CPUs **221** of each tile **210** can be coupled to a global CPU via a global CPU interface **230** (e.g., buses, connectors, sockets, etc.). This global CPU can be configured to coordinate the processing of all chiplet devices in an AI accelerator apparatus, such as apparatuses **101** and **102** of FIGS. 1A and 1B, respectively. In an example, a global CPU can use the HW DS **222** of each tile to direct each associated CPU **221** to perform various portions of the transformer computations across the slices in the tile. Also, the global CPU can be a RISC processor, or the like. The chiplet **201** also includes D2D interconnects **240** and a memory interface **250**, both of which are coupled to each of the CPUs **221** in each of the tiles. In an example, the D2D interconnects can be configured with single-ended signaling. The memory interface **250** can include one or more memory buses coupled to one or more memory devices (e.g., DRAM, SRAM, SDRAM, or the like).

(50) Further, the chiplet **201** includes a PCIe interface/bus **260** coupled to each of the CPUs **221** in each of the tiles. The PCIe interface **260** can be configured to communicate with a server or other communication system. In the case of a plurality of chiplet devices, a main bus device is coupled to the PCIe bus **260** of each chiplet device using a master chiplet device (e.g., main bus device also coupled to the master chiplet device). This master chiplet device is coupled to each other chiplet device using at least the D2D interconnects **240**. The master chiplet device and the main bus device can be configured overlying a substrate member (e.g., same substrate as chiplets or separate substrate). An apparatus integrating one or more chiplets can also be coupled to a power source (e.g., configured on-chip, configured in a system, or coupled externally) and can be configured and operable to a server, network switch, or host system using the main bus device. The server apparatus can also be one of a plurality of server apparatuses configured for a server farm within a data center, or other similar configuration.

(51) In a specific example, an AI accelerator apparatus configured for GPT-3 can incorporate eight chiplets (similar to apparatus **102** of FIG. 1B). The chiplets can be configured with D2D 16×16 Gb/s interconnects, 32-bit LPDDR5 6.4 Gb/s memory modules, and 16 lane PCIe Gen 5 PHY NRZ 32 Gb/s/lane interface. LPDDR5 (16×16 GB) can provide the necessary capacity, bandwidth and low power for large scale NLP models, such as quantized GPT-3. Of course, there can be other variations, modifications, and alternatives.

(52) FIG. 2B is a simplified block diagram illustrating an example configuration of a 16-slice chiplet device **202**. Similar to chiplet **201**, chiplet **202** includes four gangs **210** (or tiles), each of which includes four slice devices **220** and a CPU **221**. As shown, the CPU **221** of each gang/tile **210** is coupled to each of the slices **220** and to each other CPU **221** of the other gangs/tiles **210**. In an example, the tiles/gangs serve as neural cores, and the slices serve as compute cores. With this multi-core configuration, the chiplet device can be configured to take and run several computations in parallel. The CPUs **221** are also coupled to a global CPU interface **230**, D2D interconnects **240**, a memory interface **250**, and a PCIe interface **260**. As described for FIG. 2A, the global CPU interface **230** connects to a global CPU that controls all of the CPUs **221** of each gang **210**.

(53) FIG. 2C is a simplified block diagram illustrating an example configuration of a 16-slice chiplet device **203**. Chiplet **203** is similar to chiplet **201**, except that the positions of the D2D interconnects **240**, the memory interface **250**, and the PCIe interface **260** are in a different configuration. Here, a first input/output (I/O) region (shown at the top) includes one or more D2D interconnects **240** and the global CPU interface **230**, and a second I/O region (shown to the right) includes one or more D2D interconnects **240** as well. In chiplet **203**, a third I/O region (shown at the bottom) includes one or more D2D interconnects **240** and a PCIe interface **260**, whereas chiplet **201** had one or more memory interface connections **250** in this region. And, a fourth I/O region (shown to the left) includes one or more memory interface connections **250**, whereas chiplet **201** had the PCIe interface **260** in this region. In an example, these I/O regions are placed in a symmetrical configuration. The I/O placement of chiplet **203** can be used in a single die configuration for various chiplet configurations (e.g., 1×2, 2×2, 2×4, etc.). Further, the I/O placement is optimized for various array configurations due to die rotations not affecting the package I/O routing (i.e., enables scalable chiplet array configurations in any die orientation).

(54) FIG. 2D is a simplified block diagram illustrating an example configuration of a 16-slice chiplet device **204**. Similar to chiplet **202**, chiplet **204** includes four gangs **210** (or tiles), each of which includes four slice devices **220**. However, in this case, each of the slice devices **220** within each gang are coupled to a gang crossbar device **223**, which is coupled to a gang CPU and dispatch engine device **224**. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to the configurations shown in FIGS. 2A-2D.

(55) FIG. 3A is a simplified block diagram illustrating an example slice device **301** of a chiplet. For the 16-slice chiplet example, slice device **301** includes a compute core **310** having four compute paths **312**, each of which includes an input buffer (IB) device **320**, a digital in-memory-compute (DIMC) device **330**, an output buffer (OB) device **340**, and a Single Instruction, Multiple Data (SIMD) device **350** coupled together. Each of these paths **312** is coupled to a slice crossbar/controller **360**, which is controlled by the tile CPU to coordinate the computations performed by each path **312**.

(56) In an example, the DIMC is coupled to a clock and is configured within one or more portions of each of the plurality of slices of the chiplet to allow for high throughput of one or more matrix computations provided in the DIMC such that the high throughput is characterized by 512 multiply accumulates per a clock cycle. In a specific example, the clock coupled to the DIMC is a second clock derived from a first clock (e.g., chiplet clock generator, AI accelerator apparatus clock generator, etc.) configured to output a clock signal of about 0.5 GHz to 4 GHz; the second clock can be configured at an output rate of about one half of the rate of the first clock. The DIMC can also be configured to support a block structured sparsity (e.g., imposing structural constraints on weight patterns of a neural networks like a transformer).

(57) In an example, the SIMD device **350** is a SIMD processor coupled to an output of the DIMC. The SIMD **350** can be configured to process one or more non-linear operations and one or more linear operations on a vector process. The SIMD **350** can be a programmable vector unit or the like. The SIMD **350** can also include one or more random-access memory (RAM) modules, such as a data RAM module, an instruction RAM module, and the like.

(58) In an example, the slice controller **360** is coupled to all blocks of each compute path **312** and also includes a control/status register (CSR) **362** coupled to each compute path. The slice controller **360** is also coupled to a memory bank **370** and a data reshape engine (DRE) **380**. The slice controller **360** can be configured to feed data from the memory bank **370** to the blocks in each of the compute paths **312** and to coordinate these compute paths **312** by a processor interface (PIF) **364**. In a specific example, the PIF **364** is coupled to the SIMD **350** of each compute path **312**.

(59) Further details for the compute core **310** are shown in FIG. 3B. The simplified block diagram of slice device **302** includes an input buffer **320**, a DIMC matrix vector unit **330**, an output buffer **340**, a network on chip (NoC) device **342**, and a SIMD vector unit **350**. The DIMC unit **330** includes a plurality of in-memory-compute (IMC) modules **332** configured to compute a Scaled Dot-Product Attention function on input data to determine a probability distribution, which requires high-throughput matrix multiply-accumulate operations.

(60) These IMC modules **332** can also be coupled to a block floating point alignment module **334** and a partial products reduction module **336** for further processing before outputting the DIMC results to the output buffer **340**. In an example, the input buffer **320** receives input data (e.g., data vectors) from the memory bank **370** (shown in FIG. 3A) and sends the data to the IMC modules **332**. The IMC modules **332** can also receive instructions from the memory bank **370** as well.

(61) In addition to the details discussed previously, the SIMD **350** can be configured as an element-wise vector unit. The SIMD **350** can include a computation unit **352** (e.g., add, subtract, multiply, max, etc.), a look-up table (LUT) **354**, and a state machine (SM) module **356** configured to receive one or more outputs from the output buffer **340**.

(62) The NoC device **342** is coupled to the output buffer **340** configured in a feedforward loop via shortcut connection **344**. Also, the NoC device **342** is coupled to each of the slices and is configured for multicast and unicast processes. More particularly, the NoC device **342** can be configured to connect all of the slices and all of the tiles, multi-cast input activations to all of the slices/tiles, and collect the partial computations to be unicast for a specially distributed accumulation.

(63) Considering the previous eight-chiplet AI accelerator apparatus example, the input buffer can have a capacity of 64 KB with 16 banks and the output buffer can have a capacity of 128 KB with 16 banks. The DIMC can be an 8-bit block have dimensions 64×64 (eight 64×64 IMC modules) and the NoC can have a size of 512 bits. The computation block in the SIMD can be configured for 8-bit and 32-bit integer (int) and unsigned integer (uint) computations. These slice components can vary depending on which transformer the AI accelerator apparatus will serve.

(64) FIG. 4 is a simplified block diagram illustrating an example IMC module **700**. As shown, module **700** includes one or more computation tree blocks **410** that are configured to perform desired computations on input data from one or more read-write blocks **420**. Each of these read-write blocks **420** includes one or more first memory-select units **422** (also denoted as “W”), one or more second memory-select units **424** (also denoted as “I”), an activation multiplexer **426**, and an operator unit **428**. The first memory-select unit **422** provides an input to the operator unit **428**, while the second memory-select unit **424** controls the activation multiplexer **426** that is also coupled to the operator unit **428**. In the case of multiply-accumulate operations, the operator unit **428** is a multiplier unit and the computation tree blocks **410** are multiplier adder tree blocks (i.e., $\Sigma x.Math.w$).

(65) As shown in close-up **401**, each of the memory-select units **422**, **424** includes a memory cell **430** (e.g., SRAM cell, or the like) and a select multiplexer **432**. Each of the memory-select units **422**, **424** is coupled to a read-write controller **440**, which is also coupled to a memory bank/driver block **442**. In an example, the read-write controller **440** can be configured with column write drivers and column read sense amplifiers, while the memory bank/driver block **432** can be configured with sequential row select drivers.

(66) An input activation controller **450** can be coupled to the activation multiplexer **426** each of the

read-write blocks **420**. The input activation controller **450** can include precision and sparsity aware input activation register and drivers. The operator unit **428** receives the output of the first memory-select unit **422** and receives the output of this block **450** through the activation multiplexer **426**, which is controlled by the output of the second memory-select unit **424**. The output of the operator unit **428** is then fed into the computation tree block **410**.

(67) The input activation block **450** is also coupled to a clock source/generator **460**. As discussed previously, the clock generator **460** can produce a second clock derived from a first clock configured to output a clock signal of about 0.5 GHz to 4 GHz; the second clock can be configured at an output rate of about one half of the rate of the first clock. The clock generator **460** is coupled to one or more sign and precision aware accumulators **470**, which are configured to receive the output of the computation tree blocks **410**. In an example, an accumulator **470** is configured to receive the outputs of two computation tree blocks **410**. Example output readings of the IMC are shown in FIGS. **13A-13C**.

(68) Referring back to the eight-chiplet AI accelerator apparatus example, the memory cell can be a dual bank 2×6T SRAM cell, and the select multiplexer can be an 8T bank select multiplexer. In this case, the memory bank/driver block **442** includes a dual-bank SRAM bank. Also, the read/write controller can include 64 bytes of write drivers and 64 bytes of read sense amplifiers. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to these IMC module components and their configurations.

(69) FIG. **5A** is a simplified block flow diagram illustrating example numerical formats of the data being processed in a slice. Diagram **501** shows a loop with the data formats for the GM/input buffer **510**, the IMC **520**, the output buffer **530**, the SIMD **540**, and the NoC **550**, which feeds back to the GM/input buffer **510**. The IMC block **520** shows the multiply-accumulate operation ($\Sigma x.Math.w$). Additionally, the format for the data from IMC **532** flows to the output buffer **530** as well. In this example, the numerical formats include integer (int), floating point (float), and block floating (bfloat) of varying lengths.

(70) FIG. **5B** is a simplified diagram illustrating certain numerical formats, including certain formats shown in FIG. **5A**. Block floating point numerics can be used to address certain barriers to performance. Training of transformers is generally done in floating point, i.e., 32-bit float or 16-bit float, and inference is generally done in 8-bit integer (“int8”). With block floating point, an exponent is shared across a set of mantissa significant values (see diagonally line filled blocks of the int8 vectors at the bottom of FIG. **5B**), as opposed to floating point where each mantissa has a separate exponent (see 32-bit float and 16-bit float formats at the top of FIG. **5A**). The method of using block floating point numerical formats for training can exhibit the efficiency of fixed point without the problems of integer arithmetic, and can also allow for use of a smaller mantissa, e.g., 4-bit integer (“int4”) while retaining accuracy. Further, by using the block floating point format (e.g., for activation, weights, etc.) and sparsity, the inference of the training models can be accelerated for better performance. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to these numerical formats used to process transformer workloads.

(71) FIG. **6** illustrates a simplified transformer architecture **600**. The typical transformer can be described as having an encoder stack configured with a decoder stack, and each such stack can have one or more layers. Within the encoder layers **610**, a self-attention layer **612** determines contextual information while encoding input data and feeds the encoded data to a feed-forward neural network **616**. The encoder layers **610** process an input sequence from bottom to top, transforming the output into a set of attention vectors K and V. The decoder layers **620** also include a corresponding self-attention layer **622** and feed-forward neural network **626**, and can further include an encoder-decoder attention layer **624** uses the attention vectors from the encoder stack that aid the decoder in further contextual processing. The decoder stack outputs a vector of floating points (as discussed for FIG. **5B**), which is fed to linear and softmax layers **630** to project the output into a final desired result (e.g., desired word prediction, interpretation, or translation). The

linear layer is a fully-connected neural network that projects the decoder output vector into a larger vector (i.e., logits vector) that contains scores associated with all potential results (e.g., all potential words), and the softmax layer turns these scores into probabilities. Based on the this probability output, the projected word meaning may be chosen based on the highest probability or by other derived criteria depending on the application.

(72) Transformer model variations include those based on just the decoder stack (e.g., transformer language models such as GPT-2, GPT-3, etc.) and those based on just the encoder stack (e.g., masked language models such as BERT, BERT Large, etc.). Transformers are based on four parameters: sequence length (S) (i.e., number of tokens), number of attention heads (A), number of layers (L), and embedding length (H). Variations of these parameters are used to build practically all transformer-based models today. Embodiments of the present invention can be configured for any similar model types.

(73) A transformer starts as untrained and is pre-trained by exposure to a desired data set for a desired learning application. Transformer-based language models are exposed to large volumes of text (e.g., Wikipedia) to train language processing functions such as predicting the next word in a text sequence, translating the text to another language, etc. This training process involves converting the text (e.g., words or parts of words) into token IDs, evaluating the context of the tokens by a self-attention layer, and predicting the result by a feed forward neural network.

(74) The self-attention process includes (1) determining query (Q), key (K), and value (V) vectors for the embedding of each word in an input sentence, (2) calculating a score for from the dot product of Q and K for each word of the input sentence against a target word, (3) dividing the scores by the square root of the dimension of K, (4) passing the result through a softmax operation to normalize the scores, (5) multiplying each V by the softmax score, and (6) summing up the weighted V vectors to produce the output. An example self-attention process **700** is shown in FIG. 7.

(75) As shown, process **700** shows the evaluation of the sentence “the beetle drove off” at the bottom to determine the meaning of the word “beetle” (e.g., insect or automobile). The first step is to determine the $q_{\text{sub.beetle}}$, $k_{\text{sub.beetle}}$, and $v_{\text{sub.beetle}}$ vectors for the embedding vector $e_{\text{sub.beetle}}$. This is done by multiplying $e_{\text{sub.beetle}}$ by three different pre-trained weight matrices $W_{\text{sub.q}}$, $W_{\text{sub.k}}$, and $W_{\text{sub.v}}$. The second step is to calculate the dot products of $q_{\text{sub.beetle}}$ with the K vector of each word in the sentence (i.e., $k_{\text{sub.the}}$, $k_{\text{sub.beetle}}$, $k_{\text{sub.drove}}$, and $k_{\text{sub.off}}$), shown by the arrows between $q_{\text{sub.beetle}}$ and each K vector. The third step is to divide the scores by the square root of the dimension $d_{\text{sub.k}}$, and the fourth step is to normalize the scores using a softmax function, resulting in $\lambda_{\text{sub.i}}$. The fifth step is to multiply the V vectors by the softmax score ($\lambda_{\text{sub.iv.sub.i}}$) in preparation for the final step of summing up all the weight value vectors, shown by v' at the top.

(76) Process **700** only shows the self-attention process for the word “beetle”, but the self-attention process can be performed for each word in the sentence in parallel. The same steps apply for word prediction, interpretation, translation, and other inference tasks. Further details of the self-attention process in the BERT Large model are shown in FIGS. 8 and 9.

(77) A simplified block diagram of the BERT Large model (S=384, A=16, L=34, and H=1024) is shown in FIG. 8. This figure illustrates a single layer **800** of a BERT Large transformer, which includes an attention head device **810** configured with three different fully-connected (FC) matrices **821-823**. As discussed previously, the attention head **810** receives embedding inputs (384×1024 for BERT Large) and measures the probability distribution to come up with a numerical value based on the context of the surrounding words. This is done by computing different combination of softmax around a particular input value and producing a value matrix output having the attention scores.

(78) Further details of the attention head **810** are provided in FIG. 9. As shown, the attention head **900** computes a score according to an attention head function: $\text{Attention}(Q, K, V) = \text{softmax}(QK^{\text{sup.T}}/\sqrt{d_{\text{sub.k}}})V$. This function takes queries (Q), keys (K) of dimension $d_{\text{sub.k}}$,

and values (V) of dimension $d_{sub.k}$ and computes the dot products of the query with all of the keys, divides the result by a scaling factor $\sqrt{d_{sub.k}}$ and applies a softmax function to obtain the weights (i.e., probability distribution) on the values, as shown previously in FIG. 7.

(79) The function is implemented by several matrix multipliers and function blocks. An input matrix multiplier **910** obtains the Q, K, and V vectors from the embeddings. The transpose function block **920** computes $K_{sup.T}$, and a first matrix multiplier **931** computes the scaled dot product $QK_{sup.T}/\sqrt{d_{sub.k}}$. The softmax block **940** performs the softmax function on the output from the first matrix multiplier **931**, and a second matrix multiplier **932** computes the dot product of the softmax result and V.

(80) For BERT Large, 16 such independent attention heads run in parallel on 16 AI slices. These independent results are concatenated and projected once again to determine the final values. The multi-head attention approach can be used by transformers for (1) “encoder-decoder attention” layers that allow every position in the decoder to attend over all positions of the input sequence, (2) self-attention layers that allows each position in the encoder to attend to all positions in the previous encoder layer, and (3) self-attention layers that allow each position in the decoder to attend to all positions in the decoder up to and including that position. Of course, there can be variations, modifications, and alternatives in other transformer.

(81) Returning to FIG. 8, the attention score output then goes to a first FC matrix layer **821**, which is configured to process the outputs of all of the attention heads. The first FC matrix output goes to a first local response normalization (LRN) block **841** through a short-cut connection **830** that also receives the embedding inputs. The first LRN block output goes to a second FC matrix **822** and a third FC matrix **823** with a Gaussian Error Linear Unit (GELU) activation block **850** configured in between. The third FC matrix output goes to a second LRN block **842** through a second short-cut connection **832**, which also receives the output of the first LRN block **841**.

(82) Using a transformer like BERT Large, NLP requires very high compute (e.g., five orders of magnitude higher than CV). For example, BERT Large requires 5.6 giga-multiply-accumulate operations per second (“GMACs”) per transformer layer. Thus, the NLP inference challenge is to deliver this performance at the lowest energy consumption.

(83) Although the present invention is discussed in the context of a BERT Large transformer for NLP applications, those of ordinary skill in the art will recognize variations, modifications, and alternatives. The particular embodiments shown can also be adapted to other transformer-based models and other AI/machine learning applications.

(84) Many things impact the performance of such transformer architectures. The softmax function tends to be the critical path of the transformer layers (and has been difficult to accelerate in hardware). Requirements for overlapping the compute, SIMD operations and NoC transfers also impacts performance. Further, efficiency of NoC, SIMD, and memory bandwidth utilization is important as well.

(85) Different techniques can be applied in conjunction with the AI accelerator apparatus and chiplet device examples to improve performance, such as quantization, sparsity, knowledge distillation, efficient tokenization, and software optimizations. Supporting variable sequence length (i.e., not requiring padding to the highest sequence lengths) can also reduce memory requirements. Other techniques can include optimizations of how to split self-attention among slices and chips, moving layers and tensors between the slices and chips, and data movement between layers and FC matrices.

(86) According to an example, the present invention provides for an AI accelerator apparatus (such as shown in FIGS. 1A and 1B) coupled to an aggregate of transformer devices (e.g., BERT, BERT Large, GPT-2, GPT-3, or the like). In a specific example, this aggregate of transformer devices can include a plurality of transformers configured in a stack ranging from three to N layers, where N is an integer up to 128.

(87) In an example, each of the transformers is configured within one or more DIMCs such that

each of the transformers comprises a plurality of matrix multipliers including QKV matrices configured for an attention layer of a transformer followed by three fully-connected matrices (FC). In this configuration, the DIMC is configured to accelerate the transformer and further comprises a dot product of $Q \cdot K \cdot \sup.T$ followed by a softmax $(Q \cdot K \cdot \sup.T / \text{square root}(d \cdot \text{sub}.k)) \cdot V$. In an example, the AI accelerator apparatus also includes a SIMD device (as shown in FIGS. 3A and 3B) configured to accelerate a computing process of the softmax function.

(88) According to an example, the present invention provides for methods of compiling the data representations related to transformer-based models mapping them to an AI accelerator apparatus in a spatial array. These methods can use the previously discussed numerical formats as well as sparsity patterns. Using a compile algorithm, the data can be configured to a dependency graph, which the global CPU can use to map the data to the tiles and slices of the chiplets. Example mapping methods are shown in FIGS. 10-13B.

(89) FIG. 10 is a simplified table representing an example mapping process between a 24-layer transformer and an example eight-chiplet AI accelerator apparatus. As shown, the chiplets are denoted by the row numbers on the left end and the model layers mapped over time are denoted by the table entry numbers. In this case, the 24 layers of the transformer (e.g., BERT Large) are mapped to the chiplets sequentially in a staggered manner (i.e., first layer mapped onto the first chiplet, the second layer mapped onto the second chiplet one cycle after the first, the third layer mapped onto the third chiplet two cycles after the first, etc.) After eight cycles, the mapping process loops back to the first chiplet to start mapping the next eight model layers.

(90) FIG. 11 is a simplified block flow diagram illustrating a mapping process between a transformer and an example AI accelerator apparatus. As shown, a transformer 1101 includes a plurality of transformer layers 1110, each having an attention layer 1102. In this case, there are 16 attention heads 1110 (e.g., BERT Large) computing the attention function as discussed previously. These 16 attention heads are mapped to 16 slices 1130 of an AI accelerator apparatus 1103 (similar to apparatuses 201 and 202) via global CPU 1132 communicating to the slice CPUs 1134.

(91) FIG. 12 is a simplified table representing an example tiling attention process between a transformer and an example AI accelerator apparatus. Table 1200 shows positions of Q, K, and V vectors and the timing of the softmax performed on these vectors. The different instances of the softmax are distinguished by fill pattern (e.g., diagonal line filled blocks representing Q, K, V vectors and diagonal line filled blocks representing Q-K and Softmax-V dot products).

(92) In an example, the embedding E is a $[64L, 1024]$ matrix ($L=6$ for sentence length of 384), and $E \cdot \text{sub}.i$ is a $[64, 1024]$ submatrix of E, which is determined as $E \cdot \text{sub}.i = E \cdot \text{sub}.(64i-63):(64i), 1:1024$, where $i=1 \dots L$. Each of the K and Q matrices can be allocated to two slices (e.g., $@[SL1:AC3,4]: K \cdot \text{sub}.i \leftarrow E \cdot \text{sub}.i \times K \cdot \text{sub}.1 \dots 1024, 1 \dots 64$; and $@[SL1:AC1,2]: Q \cdot \text{sub}.i \leftarrow E \cdot \text{sub}.i \times Q \cdot \text{sub}.1 \dots 1024, 1 \dots 64$). An example data flows through IMC and SIMD modules are shown in the simplified tables of FIGS. 13A-13C.

(93) FIG. 13A shows table 1301 representing mapping self-attention to an AI slice according to an example of the present invention. The left side shows the IMC cycles for matrix multiplications performed by IMC modules AC1-AC4, while the right side shows SIMD cycles for element-wise computations performed by SIMD modules SIMD1-SIMD4. In this example, the IMC modules determine the key vectors K1-K6 ($a[64 \times 512]$; $w[512 \times 64]$; $o[64 \times 64]$), and query vectors Q1-Q6 ($a[64 \times 512]$; $w[512 \times 64]$; $o[64 \times 64]$), followed by the transpose QKT1-QKT6 ($a[64 \times 64]$; $w[64 \times 384]$; $o[64 \times 384]$). Then, the SIMD modules compute the softmax Smax1-Smax6 ($a[64 \times 384]$). Meanwhile, the IMC modules determine the value vectors V1-V6 ($a[64 \times 512]$; $w[512 \times 64]$; $o[64 \times 64]$), followed by the multiplication of the value vectors and the softmax results.

(94) FIG. 13B shows table 1302 representing mapping dense embedding vectors and the second FC matrix to an AI slice (left: IMCs; right: SIMDs) according to an example of the present invention. In this example, the IMCs process the embedding vectors E1-E6 ($a[64 \times 512]$; $w[512 \times 64]$; $o[64 \times 64]$), which corresponds to the path from the attention head 810 to the second FC matrix 822

in FIG. 8. Following the processing of each embedding vector E, the SIMDs process the GELU (a[64×64]), which corresponds to the path through the first LRN block **841** and the GELU block **850** in FIG. 8.

(95) FIG. 13C shows table **1303** representing mapping the third FC matrix to an AI slice (left: IMCs; right: SIMDs) according to an example of the present invention. In this example, the IMCs process the results through the second FC matrix, which corresponds to the path through the third FC matrix **823** and the second LRN block **842** in FIG. 8. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to the mappings shown in FIGS. 10-13C.

(96) According to various examples, the present invention also provides for three-dimensional (3D) stacking methods and configurations for AI accelerator apparatuses, chiplet devices, and related components. As the scope of transformer workloads expand, memory and interconnect bandwidths limit the performance of processing these workloads. Depending on the embodiment, the present 3D stacking methods and configurations can have significant advantages, such as improvement in bandwidth and power performance over conventional embodiments, reduction in number of cards required in serving systems, and low-cost fabrication processes.

(97) FIGS. 14A and 14B shows a simplified diagram illustrating a method of forming a stacked wafer apparatus according to an example of the present invention. As shown in FIG. 14A, the method **1401** starts with a memory wafer **1410** (e.g., DRAM wafer, or the like) having a plurality of memory dies (e.g., DRAM dies, or the like) formed overlying and a logic wafer **1420** having a plurality of logic dies formed overlying. Each of the logic dies can include the previously discussed AI accelerator apparatuses, chiplet devices, or related components formed overlying. Here, the memory wafer **1410** is rotated face down (shown by dotted arrow **1430**) and bonded overlying the logic wafer **1420**, which is kept in its original orientation (shown by dotted arrow **1432**). The resulting bonded wafers **1440** will have a plurality of bonded dies **1402** (only one shown for clarity) that have a memory die bonded face down overlying a logic die. Alternatively, the logic wafer **1420** can be rotated face down and bonded overlying the memory wafer **1410**. In a specific example, the method can include a three-dimensional (3D) logic-to memory hybrid bonding process. FIG. 14B shows an example bonded die more closely.

(98) As shown in FIG. 14B, the bonded die **1402** includes a memory die **1412** having a plurality of memory devices **1414** and a logic die **1422** having a plurality of logic devices **1424**. In this figure, the components are shown in an exploded view and the memory devices **1414** are visible from the topside surface region for clarity. In this configuration, the bonded die **1402** will have a plurality of bonded devices having at least a memory device **1414** bonded to a logic device **1424**. In an example, the logic device **1424** can include an AI accelerator apparatuses, a chiplet device, or related components formed overlying a first substrate. Also, the memory device **1414** can include one or more memory units (e.g., DDR DRAM devices, or the like) formed overlying a second substrate, which is mechanically and operably bonded to the first substrate and configured to a memory interface (e.g., DRAM interface, or the like) of the AI accelerator apparatus, chiplet device, etc.

(99) Or, the second semiconductor substrate can include a plurality of DRAM memory cells, one of more of the plurality of DRAM memory cells being coupled to the DRAM interface such that the first semiconductor substrate and the second semiconductor substrate are bonded through a mechanical interface. Further, a substrate member can be configured to provide mechanical support and having a surface region, the surface region being coupled to support the chiplet and memory device. In an example, these substrates can be semiconductor substrates or the like. Further, the resulting bonded device can include 3D stacked devices, such as a 3D stacked chiplet and memory device, or the like. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives.

(100) FIG. 14C is a simplified block diagram illustrating a memory device according to an example of the present invention. As shown, the memory device **1403** includes at least one or more memory

units **1450** (e.g., memory banks, or the like), row/column control units **1460**, buffer units **1462**, and routing information base (RIB) units **1464**. Here, each memory unit **1450** is coupled to and configured with a control unit **1460** and a RIB unit **1464**, and each buffer unit **1462** is coupled to and configured with four such memory units **1450**, control units **1460**, and RIB units **1464**. In a specific example, this memory device **1403** can be a DRAM core device with eight memory banks and on-chip Error Correction Code (ECC). Of course, there can be other variations, modifications, and alternatives.

(101) FIG. **15** is a simplified block illustrating a 3D stacked device die according to an example of the present invention. As shown, the stacked device die **1500** includes a plurality of logic devices **1510**, each of which includes at least a weight/activation memories module **1520** coupled to a compute engines module **1530** and a controller **1540**. In a specific example, the controller **1540** can be configured as a tiny DRAM controller. Each controller **1540** can be configured using a plurality of vias **1542** (e.g., through-silicon vias, or the like) to access a memory bank in the memory die **1550**, which is represented by the dotted line region **1550**. Although interconnects, hierarchies, and chip interfaces are not shown, this device configuration can be combined with the AI accelerator apparatus and chiplet device configurations discussed previously. For example, the weight/activations memories modules **1520** and compute engines modules **1530** can include the chiplet and slice devices and related components discussed previously. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to this stacked device configuration.

(102) FIGS. **16A** and **16B** are simplified diagrams illustrating an error correction implementation for a memory device according to an example of the present invention. As shown in FIG. **16A**, diagram **1601** shows an example memory layout, which includes information symbols (in the white blocks), outer code parity (row), inner code parity (column), and 'checks on checks'. FIG. **16B** shows a similar diagram **1602** with different configuration of dimensions but with the same memory regions for information symbols **1610**, checks on rows **1620**, checks on columns **1630**, and checks on checks **1640**. Depending on the embodiment, the memory device can implement this error correction code (ECC) on-chip.

(103) FIGS. **17A-17C** are simplified diagrams illustrating a memory tiling configuration for a stacked wafer apparatus according to an example of the present invention. FIG. **17A** shows a memory die **1701** having a plurality of memory devices **1710**, similar to the memory wafer **1410** shown in FIG. **14A**, except here a memory tile **1712** is shown outlined by the dotted line region. Here, the tile is shown in a square 2×2 configuration, but the tile can have different configurations, such as 3×3, 2×4, 4×4, etc. Using such memory tile configurations, the resulting stacked devices can have high memory bandwidth and memory capacity. Further, the scalable memory tile configurations allow for flexible composition of bandwidth and size.

(104) FIG. **17B** shows a close up of a portion of the memory devices **1710** with the previous memory tile **1712** outlined by the dotted line region. FIG. **17C** shows an exploded view of a stacked wafer apparatus **1703** in which the tile **1712** aligns to one of the logic dies **1722** on a logic wafer **1720** when bonded with the memory wafer **1701**. Here, the logic dies **1722** are shown on the topside of the logic wafer **1720** for clarity of the alignment. As discussed previously, the method of bonding can include the memory wafer **1701** rotated face down and bonded overlying the logic wafer **1720**, or vice versa.

(105) FIG. **18** is a simplified diagram illustrating a cross-sectional view of a stacked logic and memory device according to an example of the present invention. As shown, the stacked logic and memory device includes a logic device portion and a memory device portion. Here, the logic device portion includes logic Front End of Line (FEOL) layers **1810** and logic Back End of Line (BEOL) layers **1820**, while the memory device portion includes memory FEOL **1840** layers and memory BEOL layers **1850**. The logic BEOL layers **1820** include metal layers **1822**, representing interconnections between components in these layers, and metal layers **1824**, representing

interconnections from these layers. Similarly, the memory BEOL layers **1820** include metal layers **1842**, representing interconnections from these layers, and metal layers **1844**, representing interconnections between components in these layers. The logic FEOL and BEOL layers **1810**, **1820** can include implementations of the AI accelerator apparatus, chiplet device, and related components discussed previously, and the memory FEOL and BEOL layers **1840**, **1850** can include implementations of the memory devices discussed previously.

(106) These two portions can be bonded at the hybrid-bond layers **1830** via bond pads **1834** and **1836**. More specifically, the metal layers **1824** of the logic BEOL layers **1820** are coupled to a frontside redistribution layer (FRDL) contact **1832**, which is coupled to the bond pad **1834**, while the metal layers **1842** are coupled to the bond pad **1836**. The metal layers **1842** are also coupled to a through-silicon via (TSV) **1852**, which is also coupled to a memory backside redistribution layer (BRDL) contact **1862** within the packaging layer **1860**. This BRDL contact **1862** is also coupled to a bond pad **1864** on the topside of the device (i.e., the backside of the memory device). Those of ordinary skill in the art will recognize variations, modifications, and alternatives to these layer configurations.

(107) FIGS. **19A-19C** are simplified diagrams illustrating a 3D stacked chiplet and memory device using a hybrid-bonding interconnection according to various examples of the present invention. As shown in FIG. **19A**, the stacked device **1901** includes a memory die **1910** with an overlying logic die **1920**. The memory die **1910** includes top metal (TM) contacts **1912** and bonding interconnection materials **1914** within dielectric layers **1930**, while the logic die **1920** includes TM contacts **1922** and bonding interconnection materials **1924** within dielectric layers **1932**. The logic die **1920** is coupled to the memory die **1910** at the bonding interface **1940** such that the bonding interconnection materials **1914** are bonded to the bonding interconnection materials **1924**.

(108) The bonding interconnection materials **1914** and **1924** can include metal materials, such as Cu, Al, and the like. In a specific example, the bonding method can include a direct fusion process (e.g., Cu—Cu direct fusion, or the like) with a low bonding temperature (e.g., less than 350 degrees Celsius). In FIG. **19A**, an example chemical composition within the dielectric layers **1930** and **1932** is shown as having a silicon oxide based composition, but other dielectric materials may be used as well.

(109) FIG. **19B** shows another stacked device **1902** but with a pair of bonding interconnection materials **1914** coupled to one TM contact **1912** and a pair of bonding interconnection materials **1924** coupled to one TM contact **1922**, which leads to a multi-point bond at the bonding interface **1940**. Device **1902** also includes another bonding interconnection in which the TM metal **1922** is also coupled to a metal layer **1940** which is coupled to a bonding structure **1942** that is exposed via an opening in the logic die **1920**.

(110) FIG. **19C** shows another stacked device **1903** which has similar TM contacts **1912**, **1922** and bonding interconnection materials **1914**, **1924**, but also includes align marker materials **1916**, **1926**. The align marker materials **1916** are aligned to the TM contacts **1912** and the align marker materials **1926** are aligned to the TM contacts **1922**. In a specific example, the aligned markers can be formed with a high precision (e.g., 0.25 μm). The stacked device can also have small pitch size (e.g., 3 μm) and up to 110,000/mm² integration density. Depending on the embodiment, such stacked devices can include implementations of the AI accelerator apparatus, chiplet device, memory devices, and related components discussed previously.

(111) FIG. **19D** is a device level image showing a cross-sectional view of a 3D stacked chiplet and memory device according to an example of the present invention. As shown, device image **1904** shows a logic TM contact coupled to a DRAM TM contact via bonding interconnection materials coupled at the bonding interface **1940**. Those of ordinary skill in the art will recognize variations, modifications, and alternatives to these bonding techniques.

(112) FIG. **20** is a simplified diagram illustrating a top view of a 3D stacked AI accelerator apparatus according to an example of the present invention. As shown, the apparatus **2000** includes

a plurality of 3D stacked chiplet devices **2030** formed overlying a wafer substrate **2010** with a die region **2020**. As discussed previously, the stacked chiplet devices **2030** can be formed using a wafer-on-wafer hybrid bonding process, or the like. Here, there are four stacked chiplet devices **2030** arranged in a 2×2 configuration overlying the die region **2020**, but there can be other configurations (e.g., 1×2, 3×3, 2×4, etc.). Further details of these chiplet devices **2030** are shown in FIGS. **21A-21F**.

(113) FIG. **21A** is a simplified diagram illustrating a cross-sectional view of an example 3D stacked chiplet device **2101** with a logic die **2110** overlying a memory die **2120**. These dies are bonded by a plurality of contacts **2130**, which can include the materials and processes discussed previously. In a specific example, the memory die **2120** can be formed as a thin layer (e.g., 10 um memory die vs. 775 um logic die for 21 nm technology) that extends as if part of the logic die **2110**. Also, the plurality of contacts **2130** can be characterized by 3 u pitch and the bump contacts **2128** can be characterized by a 110 um-130 um pitch. Here, the memory die **2120** shows BEOL layers **2122** and FEOL layers **2124** with TSVs **2126** configured between these layers and coupled to bump contacts **2128**.

(114) FIG. **21B** is a simplified diagram illustrating a cross-sectional view of an example 3D stacked chiplet device **2102** with the memory die **2120** overlying the logic die **2110**. These dies are also bonded by the plurality of contacts **2130**. Here, the logic die **2110** shows BEOL layers **2112** and FEOL layers **2114** with TSVs **2116** configured between these layers and coupled to bump contacts **2118**.

(115) FIG. **21C** is a simplified diagram illustrating a cross-sectional view of an example 3D stacked chiplet device **2103** with stacked memory dies **2020**, **2040** overlying the logic die **2110**. Similar to device **2102**, these dies are bonded by the plurality of contacts **2130**. Here, the second memory die **2140** includes BEOL layers **2142** and FEOL layers **2144**, and is bonded to the first memory die **2120** using a face-to-back hybrid bonding process. This face-to-back bond results in the FEOL layers **2124** of the first memory die **2120** being bonded to the BEOL layers **2142** of the second memory die **2140**. Depending on the embodiment, additional memory dies can be stacked in the same manner for higher capacity. Further, the logic die **2010** can be configured overlying the stacked memory dies (similar to device **2101**) as well.

(116) FIGS. **21D** and **21E** are simplified diagrams illustrating cross-sectional views of example 3D stacked chiplet devices **2104** and **2105** using micro bumps **2132**. As shown, devices **2104** and **2105** are similar to devices **2102** (memory die on logic die) and **2101** (logic die on memory die), respectively, but the logic die **2110** and the memory die **2120** are bonded using micro bumps **2132**. In a specific example, the micro bumps **2132** can be characterized by a 10 u-36 um pitch. These micro bumps **2132** can also be used to form stacked memory dies, similar to device **2103**.

(117) FIG. **21F** is a simplified diagram illustrating a cross-sectional view of an example 3D stacked die-to-die (D2D) link system. As shown, system **2106** includes two 3D stacked chiplet devices similar to device **2101** overlying a substrate member **2140** (e.g., organic substrate, or the like). These two stacked chiplet devices are coupled together by a D2D interconnect **2150** coupled between the bumps **2128** of each stacked chiplet device. In an example, the interconnect **2150** can include Universal Chiplet Interconnect Express (UCIe), or the like. Those of ordinary skill in the art will recognize variations, modifications, and alternatives to these stacked configurations.

(118) FIG. **22** is a simplified diagram illustrating a top view of an example AI accelerator apparatus using full mesh chiplet connectivity according to an example of the present invention. As shown, the apparatus **2200** includes a plurality of chiplet devices **2210**, which can include 3D stacked chiplet device configurations discussed previously. Each chiplet device **2210** includes at least a CPU **2220**, a plurality of D2D interconnects **2230**, and a plurality of device interconnects **2240** (e.g., PCIe interconnects, or the like). Each chiplet device **2210** also includes a plurality of non-diagonal links both within the chiplet device **2210** and between chiplet devices **2210**. Here, apparatus **2200** shows links between the D2D interconnects **2230** within (intra-chiplet non-diagonal

links) and between (inter-chiplet non-diagonal links) the chiplet devices **2210**, and these links can be configured as active links **2232**, disabled links **2234**, or unused links **2236**. The active links **2232** provide connections between chiplet devices **2210** and to different portions of a chiplet device **2210**, while the disabled links **2234** can limit undesired connections that were provided by default. And, the unused links **2236** can be used for future connections to other components or devices (e.g., co-package optical link). In this configuration, the apparatus **2200** can have full mesh chiplet connectivity without diagonal links on the package. Of course, there can be other variations, modifications, and alternatives.

(119) These techniques can be implemented with any of the other AI accelerator apparatus configurations discussed herein. For example, the apparatus can include a plurality of chiplets coupled together in the full mesh connectivity configuration. Each of these chiplets can include a plurality of tiles, and each of these tiles can include at least a plurality of slices, a CPU coupled to the plurality of slices, and a hardware dispatch device coupled to the CPU. The apparatus includes a plurality of D2D interconnects coupled to each of the CPUs in each of the tiles, and at least one of the D2D interconnects of each chiplet is coupled to one of the D2D interconnects of each other chiplet using a non-diagonal link. In a specific example, the plurality of chiplets includes at least four chiplets, and each of the chiplets is coupled to each other chiplet in the full mesh connectivity configuration using a plurality of intra-chiplet non-diagonal links and a plurality of inter-chiplet non-diagonal links.

(120) The apparatus can also include a first clock configured to output a clock signal of about 0.5 GHz to 4 GHz, and each of the slices can include a digital in memory compute (DIMC) device coupled to a second clock configured at an output rate of one half of the rate of the first clock. As discussed previously, this DIMC device can be configured to allow for a throughput of one or more matrix computations provided in the DIMC device such that the throughput is characterized by 512 multiply accumulates per a clock cycle. The chiplets can also include interconnect interfaces (e.g., PCIe interfaces, or the like), memory interfaces (e.g., DRAM interfaces, or the like), global CPU interfaces (e.g., RISC interfaces, or the like), as well as other interfaces to facilitate communication between the chiplets, memory and a server or host system.

(121) According to an example, the present invention provides for methods and devices for processing transformer workloads involving prompt processing. The prompt processing can include AI-driven processing tools (e.g., ChatGPT, or the like) that generate tokens in response to a given prompt, and then generate tokens in response to a follow-up prompt with context. Such prompt processing can include matrix processing using a language model, such as the BERT Large model **800** shown in FIG. **8**, with attention head devices, such as the attention head device **900** shown in FIG. **9**.

(122) For example, consider an initial prompt of “who manufactures the chips used in the Apple iPhone?”, which generates output tokens including a list of suppliers and the statement “however, in recent years, TSMC has become the primary supplier of the chips in the Apple iPhone.” A follow-up prompt can include “why has TSMC become the primary supplier?”, which would trigger token generation to the follow-up prompt with context. An example processing of the follow-up prompt is shown in the following figure.

(123) FIG. **23** is a simplified flow diagram illustrating a method of operating a transformer model according to an example of the present invention. The transformer model **2300** is represented by a plurality of layers **2310** configured to cyclically process input tokens using an attention mechanism. Each layer of this model includes at least the following steps/layers: an embedding layer **2320**, a Query, Key, Value (QKV) projection layer **2322**; first matrix computation layer **2324**, a softmax layer **2326**, a second matrix computation layer **2328**, a feed-forward network (FFN) layer **2330**, and a language model (LM) head layer **2332**. In an example, the layers marked by dotted line region **2312** include the self-attention steps discussed previously.

(124) In the embedding layer **2320**, the input prompt processed as input vectors (i.e., embeddings)

of each unit of the prompt. Referring to the previous example, this transformer model **2300** is shown to process the follow-up prompt “why has TSMC become the primary supplier”. As shown, each cycle through the plurality of layers **2310** is processing one word of the prompt, and the embedding layer **2320** is generating an embedding vector of each such word.

(125) In the QKV projection layer **2322**, three vectors (query, key, and value vectors) are created for each input token (e.g., embedding vector). These vectors are determined by multiplying the input token by three weight matrices trained during the training process. In a specific example, this layer **2322** implements a general matrix multiply (GeMM) algorithm to process the prompt. Also, token generation can include a general matrix vector (GeMV) algorithm. These processes can face challenges of being compute bound and/or memory bound.

(126) In the first matrix computation layer **2324**, a score is determined between a target token and each token in the sequence. This score is calculated by the dot product of the target token's query (Q) vector and the each token's key (K) vector (in a transposed format). In a specific example, the scores are also divided by the square root on the dimension (d.sub.k) of the key vectors. In the softmax layer **2326**, a softmax operation is performed to normalize the scaled scores. Then, in the second matrix computation layer **2328**, each value vector is multiplied by the softmax score to determine weighted value vectors, and the weighted value vectors are added to produce the output for the target token.

(127) These layers, which are marked by the dotted line region **2314**, can include a caching process (e.g., large KV cache, or the like) to facilitate the computations discussed previously. As shown, FIG. **23** shows example growth of QKV data stored in main memory/cache memory with each subsequent token processing cycle. In an example, these operations are condensed using matrix operations. The input tokens can be configured in a matrix and multiplied with the three weight matrices to produce query, key, and value matrices. The scores are determined by the dot product of the query and transposed key matrices, and the softmax of those scores after scaling is multiplied with the value matrix to determine an output matrix.

(128) The outputs of the self-attention layers are sent to the FFN layer **2330**, which then outputs to the LM head layer **2332**. This LM head layer **2332** can predict the next token in a sequence based on the previous tokens, which is fed to the next layer. Here, the first layer **2310** processes the “TSMC” token, the subsequent layer processes the “has” token with context, and the layer after that processes the “become” token. This cycle can process subsequent tokens with context from the previously processed token, and this cycle can continue until all tokens from the prompt are processed. Of course, those of ordinary skill in the art will recognize other variations, modifications, and alternatives.

(129) FIG. **24** is a simplified graph illustrating the effect of memory bandwidth and interconnect bandwidth on performance of processing transformer workloads. As shown, graph **2400** shows normalized scaling of hardware (HW) flops, memory bandwidth (BW), and interconnect BW over time (years). The HW flops scaling is shown to have increased about 90,000 times over 20 years (about 3.1 times over 2 years), while the memory (DRAM) BW and interconnect BW scaling is shown to have increased about 30 times over 20 years (about 1.4 times over 2 years). Further, various data points show specific reference devices for context.

(130) FIG. **25** is a simplified graph illustrating the evolution of byte/flop (B/F) ratio for deep neural networks. As shown, graph **2500** shows byte/flop ratio over time (years), with the ratio growing steadily through 2019. Further, each data point includes example neural networks for the specific year in which the byte/flop ratio is measured. At 0.06 B/F, a memory bandwidth of 25 TB/s is needed. However, AI compute applications in 2020 involve about 400 tera operations per second (TOPS). For example, GPT has a 0.125 B/F ratio, which requires about 200 TB/s. Such data shows the insatiable memory bandwidth required by AI compute applications using deep neural networks.

(131) FIG. **26** is a simplified graph illustrating the power efficiency of high bandwidth memory (HBM). As shown, graph **2600** shows estimated power (W) over bandwidth (TB/s) for two

versions of HBM: HBM-2e and HBM-3. This data shows the energy efficiency gap is getting bigger. In an example, desirable target power efficiency for AI compute can include a memory bandwidth of 25 TB/s with power at 100 W. Referring to the previous GPT example, HBM at 200 TB/s would require about 6 KW, which is too much power consumed.

(132) As discussed previously, the present methods and configurations for 3D stacked devices offer significant performance advantages. Depending on the embodiment, the benefits can include around 10 to 20 times improvement in bandwidth and power over conventional embodiments, such as HBM and the like. Stacked DRAM configurations can have around six times the capacity compared to SRAM or similar configurations. And, the capacity of such stacked configurations can be expanded through multi-layer memory device stacking. Further, the methods for fabricating these stacked configurations can have lower cost due to factors such as not requiring a silicon interposer and using proven low-cost hybrid bonding processes that are foundry compatible.

(133) FIG. 27A is a simplified graph showing throughput data for 3D stacked AI accelerator apparatuses according to various examples of the present invention. More specifically, graph 2701 shows the throughput per card for two 3D stacked apparatus configurations, “3D DRAM” and “3D DRAM 2×ACs”, as applied to a variety of transformer models. These apparatus configurations can be configured in various card form factors, such as the PCIe card examples shown in FIGS. 1A and 1B. Also, these throughput measurements are normalized to a device configuration using HBM. In this case, the data shows about a two to five times improved throughput per card for both stacked configurations compared to an HBM configuration at extremely large context lengths.

(134) FIG. 27B is a simplified graph showing latency data for 3D stacked AI accelerator apparatuses according to various examples of the present invention. More specifically, graph 2702 shows latency per token for two 3D stacked apparatus configurations, “3D DRAM” and “3D DRAM 2×ACs”, and an HBM device configuration, as applied to a variety of transformer models. In this case, the data shows that the stacked configurations have about a three to four times competitive advantage for large models, such as PALM and GPT4.

(135) FIG. 27C is a simplified graph showing a comparison of throughput per card, latency per token, and number of cards in the system for an HBM device configuration and two 3D stacked apparatus configurations according to examples of the present invention. The two stacked configurations include “3D DRAM” and “3D DRAM 2×ACs”, the same configurations shown previously in FIGS. 27A and 27B. Here, graph 2703 shows that the stacked configurations not only demonstrate the improved throughput discussed previously, but also these stacked configurations also enable around a three times reduction in number of cards in the serving system.

(136) FIG. 28A is a simplified block diagram illustrating a 3D stacked AI accelerator apparatus with a memory die having memory bank group devices according to an example of the present invention. As shown, apparatus 2801 includes a logic die 2810 coupled to a memory die 2820 and having a plurality of compute paths that extend across both dies. In the logic die 2810, the compute path includes an activation memory device 2830 coupled to compute device 2840 that is coupled to a crossbar device 2850, which is also coupled to plurality of memory controls 2860. The activation memory device 2830 includes a plurality of activation memory units that are each coupled to one of a plurality of compute units in the compute device 2840 (shown by arrows). Here, the activations are stored on the logic die 2810 (e.g., AI accelerator apparatus or chiplet device), while the weights and KV cache (i.e., cache memory devices configured for the Key and Value computations discussed previously) are offloaded to the memory die 2820. Each of these compute units and each of the memory controllers 2860 is also coupled to the crossbar device 2850 (also shown by arrows). The logic die 2810 can incorporate any of the previously discussed components and techniques discussed previously for AI accelerator apparatuses and chiplet devices.

(137) In the memory die 2820, the compute path includes a plurality of memory bank group devices 2870, each of which is also coupled to the memory controllers 2860 via interconnections 2822. Further details of these bank group devices 2870 are discussed with reference to FIG. 28C. In

an example, the interconnections **2822** can include bump connections (e.g., microbumps, or the like) or a hybrid bond connection. The interconnections **2822** can include any of the stacked bonding configurations discussed previously. This 3D stacked configuration can be optimized for large read transfers to facilitate the processing of transformer workloads.

(138) FIG. **28B** is a simplified block diagram illustrating a 3D stacked AI accelerator apparatus with multiple stacked memory dies having memory bank group devices according to an example of the present invention. As shown, apparatus **2802** extends the 3D stacking configuration of apparatus **2801** with one or more additional memory dies **2820**. In an example, the connections from the plurality of memory controls **2860** through any of the memory dies **2820** can include via structures **2824** that couple the plurality of memory bank group devices **2870** of one memory die **2820** to those of an adjacent memory die **2820**.

(139) FIG. **28C** is a simplified block diagram illustrating a memory bank group device configuration according to an example of the present invention. As shown, the device configuration **2803** includes a memory bank group device **2872** with a plurality of memory bank units **2874** (shown numbered 1 to N) and a plurality of error correction code (ECC) bank units **2876** (shown numbered 1 to M). While the memory bank units **2874** store the data for processing transformer workloads, the ECC bank units provide error correction capability to tolerate memory bank failures and can protect against microbump (ubump) manufacturing failures. For example, the bank group device **2872** configured for 64-bit bank units can have eight bank units **2874** and two ECC bank units **2876**, which can tolerate up to two bank failures. This capability can also include determining which bit locations have failed. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to the memory bank group device and 3D stacked AI accelerator apparatus configurations.

(140) FIG. **29A** is a simplified diagram illustrating a cross-sectional view of a 3D stacked AI accelerator apparatus with an organic substrate according to an example of the present invention. As shown, the apparatus **2901** includes an organic substrate **2910** with an overlying memory die **2920** and an overlying logic die **2930** in a stacked configuration. Here, the logic die **2930** is configured overlying the memory die **2920**, but the memory die **2920** can be configured overlying the logic die **2930** as well (see FIGS. **21A-21F**).

(141) The organic substrate **2910** includes a substrate core layer **2912** having a plurality of passive embedded decoupling capacitors **2914**. The organic substrate **2910** can also include a plurality of underlying packaging interconnections **2916**, which can include ball grid array (BGA) balls, or the like. The memory die **2920** includes a plurality of active embedded decoupling capacitors **2922** and a plurality of via structures **2924** (e.g., through-silicon vias, or the like). The memory die **2920** can also include any of the memory die components and configurations discussed previously. The logic die **2930** also includes a plurality of active embedded decoupling capacitors. Similarly, the logic die **2930** can also include any of the logic die components and configurations discussed previously. In a specific example, the logic die can have a thickness of about 750-800 μm and the memory die can have a thickness of about 45-55 μm , but these thicknesses can vary depending on the application.

(142) FIG. **29A** also shows an example interconnection configuration between the organic substrate **2910**, the memory die **2920**, and the logic die **2930**. Here, the memory die **2920** is coupled to the organic substrate **2910** via an underfill layer **2940** having a plurality of bump connections **2942**, which include C4 bumps (e.g., 110-130 μm pitch), or the like. These bump connections **2942** are coupled to the via structures **2924**. Similarly, the logic die **2930** is coupled to the memory die **2920** via another underfill layer **2950** having a plurality of bump connections **2952**, which can include ubumps (e.g., 30-40 μm pitch), or the like. Further, a molding compound **2954** can be around the sides of the logic die **2930** and overlying any exposed portions of the underfill layer **2950** and exposed portions of the underlying memory die **2920**.

(143) FIG. **29B** is a simplified diagram illustrating a cross-sectional view of a 3D stacked AI accelerator apparatus with an organic substrate and an embedded capacitor die according to an

example of the present invention. As shown, apparatus **2902** is similar to the previous apparatus **2901** except that the logic die **2930** does not include the plurality of active embedded decoupling capacitors **2932**. Here, the apparatus **2902** includes an embedded capacitor die **2960**, which can include a die with embedded deep trench capacitor (eDTC), and this capacitor die **2960** is configured overlying the logic die **2930**. This die **2960** is coupled to the logic die **2930** via bond interconnection **2962**, which can include a wafer-on-wafer (WoW) hybrid bond (e.g., about 2-3 um pitch, or the like).

(144) FIG. **29C** is a simplified diagram illustrating a cross-sectional view of a 3D stacked AI accelerator apparatus with an organic substrate and multiple stacked memory dies according to an example of the present invention. As shown, apparatus **2903** extends the 3D stacking configuration of apparatus **2901** with one or more additional memory dies **2920**. These additional memory dies **2920** are also coupled together by underfill layers **2950** with pluralities of bump connections **2962**. Those of ordinary skill in the art will recognize other variations, modifications, and alternatives to the stacked configurations shown in FIGS. **29A-29C**.

(145) While the above is a full description of the specific embodiments, various modifications, alternative constructions and equivalents may be used. As an example, the AI accelerator apparatus and chiplet devices can include any combination of elements described above, as well as outside of the present specification. Therefore, the above description and illustrations should not be taken as limiting the scope of the present invention which is defined by the appended claims.

Claims

1. An AI accelerator apparatus configured with in-memory compute, the apparatus comprising: a plurality of chiplets coupled together in a full mesh connectivity configuration, each of the chiplets comprising a plurality of tiles, and each of the tiles comprising: a plurality of slices, a central processing unit (CPU) coupled to the plurality of slices, and a hardware dispatch device coupled to the CPU; a first clock configured to output a clock signal of 0.5 GHz to 4 GHz; a plurality of die-to-die (D2D) interconnects coupled to the each of CPUs in each of the tiles, wherein at least one of the D2D interconnects of each chiplet is coupled to one of the D2D interconnects of each other chiplet using a non-diagonal link; a peripheral component interconnect express (PCIe) bus coupled to the CPUs in each of the tiles; a dynamic random access memory (DRAM) interface coupled to the CPUs in each of the tiles; a global reduced instruction set computer (RISC) interface coupled to each of the CPUs in each of the tiles; wherein each of the slices includes a digital in memory compute (DIMC) device coupled to a second clock and configured to allow for a throughput of one or more matrix computations provided in the DIMC device such that the throughput is characterized by 512 multiply accumulates per a clock cycle; wherein the DIMC device is coupled to the second clock configured at an output rate of one half of the rate of the first clock; and a substrate member configured to provide mechanical support and having a surface region, the surface region being coupled to support the plurality of chiplets.
2. The apparatus of claim 1 further comprising one or more double data rate (DDR) DRAM devices, the one or more DDR DRAM devices being coupled to one or more chiplets using the DRAM interface.
3. The apparatus of claim 1 further comprising a main bus device, the main bus device being coupled to each PCIe bus in each chiplet using a master chiplet device, the master chiplet device being coupled to each of the other chiplet devices using at least the plurality of D2D interconnects.
4. The apparatus of claim 3 further comprising a server apparatus, the apparatus being configured and operable to the server apparatus using the main bus device.
5. The apparatus of claim 4 wherein the server apparatus is one of a plurality of server apparatuses configured for a server farm within a data center.
6. The apparatus of claim 5 further comprising a power source coupled to the apparatus.

7. The apparatus of claim 1 further comprising an aggregate of transformer devices, the transformer devices comprising a plurality of transformers each of which is stacked in a layer by layer ranging from three (3) to M, where M is an integer up to 128.
8. The apparatus of claim 7 wherein each of the plurality of transformers is configured within one or more DIMC devices such that each of the transformers comprises a plurality of matrix multipliers including a query key value (QKV) matrices configured for an attention layer of a transformer followed by three fully connected (FC) matrices.
9. The apparatus of claim 8 wherein the DIMC device is configured to accelerate the transformer and further comprises a dot product of $QK \cdot \text{sup.T}$ followed by a softmax $(QK \cdot \text{sup.T} / \text{square root}(d \cdot \text{sub.k}))V$.
10. The apparatus of claim 9 wherein each of the slices includes a single input multiple data (SIMD) device configured to accelerate a computing process of the softmax.
11. The apparatus of claim 1 wherein each of the chiplets comprises four tiles arranged symmetrical to each other, each of the tiles comprises four slices.
12. The apparatus of claim 1 wherein the DIMC device is configured to support one or more block floating point data types using a shared exponent.
13. The apparatus of claim 12 wherein the DIMC device is configured to support a block structured sparsity.
14. The apparatus of claim 1 further comprising a network on chip (NoC) device configured for a multicast process and coupled to each of the plurality of slices.
15. The apparatus of claim 1 wherein the plurality of chiplets are configured to process a workload of a transformer; wherein the transformer includes a plurality of transformer layers, each of the transformer layers having an attention layer associated with a portion of the workload; and wherein each attention layer is mapped on to one of the plurality of slices using the global RISC interface to communicate with the CPU associated with the tile of the slice to process the portion of the workload associated with the attention layer.
16. The apparatus of claim 1 wherein the substrate member includes an interposer, and wherein the plurality of chiplets is coupled to each other using the interposer.
17. The apparatus of claim 1 further comprising a first semiconductor substrate including the plurality of chiplets; and a second semiconductor substrate comprising a plurality of DRAM memory cells, one of more of the plurality of DRAM memory cells being coupled to the DRAM interface such that the first semiconductor substrate and the second semiconductor substrate are bonded through a mechanical interface.
18. The apparatus of claim 1 wherein the plurality of D2D interconnects comprises a plurality of universal chiplet interconnect express (UCIe) interconnects.
19. The apparatus of claim 1 wherein the plurality of D2D interconnects of each chiplet are configured in at least four input/output (I/O) regions in a symmetrical configuration enabling scalable chiplet array configurations in any die orientation.
20. An AI accelerator apparatus configured with in-memory compute, the apparatus comprising: at least four chiplets, each of the chiplets comprising a plurality of tiles, and each of the tiles comprising: a plurality of slices, a central processing unit (CPU) coupled to the plurality of slices, and a hardware dispatch device coupled to the CPU; a first clock configured to output a clock signal of 0.5 GHz to 4 GHz; a plurality of die-to-die (D2D) interconnects coupled to the each of CPUs in each of the tiles, wherein each of the chiplets is coupled to each other chiplet in a full mesh connectivity configuration using a plurality of intra-chiplet non-diagonal links and a plurality of inter-chiplet non-diagonal links; a peripheral component interconnect express (PCIe) bus coupled to the CPUs in each of the tiles; a dynamic random access memory (DRAM) interface coupled to the CPUs in each of the tiles; a global reduced instruction set computer (RISC) interface coupled to each of the CPUs in each of the tiles; wherein each of the slices includes a digital in memory compute (DIMC) device coupled to a second clock and configured to allow for a

throughput of one or more matrix computations provided in the DIMC device such that the throughput is characterized by 512 multiply accumulates per a clock cycle; wherein the DIMC device is coupled to the second clock configured at an output rate of one half of the rate of the first clock; and a substrate member configured to provide mechanical support and having a surface region, the surface region being coupled to support the plurality of chiplets.

21. The apparatus of claim 20 further comprising one or more double data rate (DDR) DRAM devices, the one or more DDR DRAM devices being coupled to one or more chiplets using the DRAM interface.

22. The apparatus of claim 20 further comprising a main bus device, the main bus device being coupled to each PCIe bus in each chiplet using a master chiplet device, the master chiplet device being coupled to each of the other chiplet devices using at least the plurality of D2D interconnects.

23. The apparatus of claim 22 further comprising a server apparatus, the apparatus being configured and operable to the server apparatus using the main bus device.

24. The apparatus of claim 23 wherein the server apparatus is one of a plurality of server apparatuses configured for a server farm within a data center.

25. The apparatus of claim 24 further comprising a power source coupled to the apparatus.

26. The apparatus of claim 20 further comprising an aggregate of transformer devices, the transformer devices comprising a plurality of transformers each of which is stacked in a layer by layer ranging from three (3) to M, where M is an integer up to 128.

27. The apparatus of claim 26 wherein each of the plurality of transformers is configured within one or more DIMC devices such that each of the transformers comprises a plurality of matrix multipliers including a query key value (QKV) matrices configured for an attention layer of a transformer followed by three fully connected (FC) matrices.

28. The apparatus of claim 27 wherein the DIMC device is configured to accelerate the transformer and further comprises a dot product of $QK \cdot \text{sup.T}$ followed by a softmax $(QK \cdot \text{sup.T} / \text{square root}(d \cdot \text{sub.k}))V$.

29. The apparatus of claim 28 wherein each of the slices includes a single input multiple data (SIMD) device configured to accelerate a computing process of the softmax.

30. The apparatus of claim 20 wherein each of the chiplets comprises four tiles arranged symmetrical to each other, each of the tiles comprises four slices.

31. The apparatus of claim 20 wherein the DIMC device is configured to support one or more block floating point data types using a shared exponent.

32. The apparatus of claim 20 wherein the DIMC device is configured to support a block structured sparsity.

33. The apparatus of claim 20 further comprising a network on chip (NoC) device configured for a multicast process and coupled to each of the plurality of slices.

34. The apparatus of claim 20 wherein the chiplets are configured to process a workload of a transformer; wherein the transformer includes a plurality of transformer layers, each of the transformer layers having an attention layer associated with a portion of the workload; and wherein each attention layer is mapped on to one of the plurality of slices using the global RISC interface to communicate with the CPU associated with the tile of the slice to process the portion of the workload associated with the attention layer.

35. The apparatus of claim 20 wherein the substrate member includes an interposer, and wherein the chiplets are coupled to each other using the interposer.

36. The apparatus of claim 20 further comprising a first semiconductor substrate including the chiplets; and a second semiconductor substrate comprising a plurality of DRAM memory cells, one of more of the plurality of DRAM memory cells being coupled to the DRAM interface such that the first semiconductor substrate and the second semiconductor substrate are bonded through a mechanical interface.

37. The apparatus of claim 20 wherein the plurality of D2D interconnects comprises a plurality of

universal chiplet interconnect express (UCIe) interconnects.

38. The apparatus of claim 20 wherein the plurality of D2D interconnects of each chiplet are configured in at least four input/output (I/O) regions in a symmetrical configuration enabling scalable chiplet array configurations in any die orientation.

39. An AI accelerator apparatus configured with in-memory compute, the apparatus comprising: a plurality of chiplets coupled together in a full mesh connectivity configuration, each of the chiplets comprising a plurality of tiles, and each of the tiles comprising: a plurality of slices, and a central processing unit (CPU) coupled to the plurality of slices; a plurality of die-to-die (D2D) interconnects coupled to the each of CPUs in each of the tiles, wherein at least one of the D2D interconnects of each chiplet is coupled to one of the D2D interconnects of each other chiplet using a non-diagonal link; wherein each of the slices includes a digital in memory compute (DIMC) device; and a substrate member configured to provide mechanical support and having a surface region, the surface region being coupled to support the plurality of chiplets.

40. The apparatus of claim 39 wherein the plurality of chiplets includes at least four chiplets; and wherein each of the chiplets is coupled to each other chiplet in a full mesh connectivity configuration using a plurality of intra-chiplet non-diagonal links and a plurality of inter-chiplet non-diagonal links.
