

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386827
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Ramanasankaran; Rajiv et al.

Building data platform with a distributed digital twin

Abstract

A method including receiving, by one or more processing circuits, building data, generating, by the one or more processing circuits, a first digital twin based on the building data, wherein a first system stores the first digital twin and a second system stores a second digital twin generated based on the building data, where the first digital twin includes a relationship that forms a connection between the first digital twin and the second digital twin by linking a first entity of the first entities of the first digital twin and a second entity of the second entities of the second digital twin, and performing, by the one or more processing circuits, one or more operations based on at least one of the first digital twin, the second digital twin, or the relationship that forms the connection between the first digital twin and the second digital twin.

Inventors: Ramanasankaran; Rajiv (San Jose, CA), Shatdal; Ambuj (Madison, WI), O'Dierno; Dominick James (Mt Pleasant, WI), Otto; Brian Scott (Menomonee Falls, WI), Galvez; Miguel (Westford, MA), Swanson; Trent Mark (Wellington, FL)

Applicant: Tyco Fire & Security GmbH (Schaffhausen, CH)

Family ID: 86383908

Assignee: TYCO FIRE & SECURITY GMBH (Neuhausen am Rheinfall, CH)

Appl. No.: 18/143968

Filed: May 05, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230273922 A1	Aug. 31, 2023

Related U.S. Application Data

continuation parent-doc US 17534821 20211124 US 11704311 child-doc US 18143968

Publication Classification

Int. Cl.: G06F15/16 (20060101); **G06F16/21** (20190101); **G06F16/2453** (20190101); **G06F16/2455** (20190101); **G06F16/27** (20190101)

U.S. Cl.:

CPC G06F16/24532 (20190101); **G06F16/211** (20190101); **G06F16/24553** (20190101); **G06F16/275** (20190101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5301109	12/1993	Landauer et al.	N/A	N/A
5446677	12/1994	Jensen et al.	N/A	N/A
5581478	12/1995	Cruse et al.	N/A	N/A
5812962	12/1997	Kovac	N/A	N/A
5960381	12/1998	Singers et al.	N/A	N/A
5973662	12/1998	Singers et al.	N/A	N/A
6014612	12/1999	Larson et al.	N/A	N/A
6031547	12/1999	Kennedy	N/A	N/A
6134511	12/1999	Subbarao	N/A	N/A
6157943	12/1999	Meyer	N/A	N/A
6285966	12/2000	Brown et al.	N/A	N/A
6363422	12/2001	Hunter et al.	N/A	N/A
6385510	12/2001	Hoog et al.	N/A	N/A
6389331	12/2001	Jensen et al.	N/A	N/A
6401027	12/2001	Xu et al.	N/A	N/A
6437691	12/2001	Sandelman et al.	N/A	N/A
6477518	12/2001	Li et al.	N/A	N/A
6487457	12/2001	Hull et al.	N/A	N/A
6493755	12/2001	Hansen et al.	N/A	N/A
6577323	12/2002	Jamieson et al.	N/A	N/A
6626366	12/2002	Kayahara et al.	N/A	N/A
6646660	12/2002	Patty	N/A	N/A
6704016	12/2003	Oliver et al.	N/A	N/A
6732540	12/2003	Sugihara et al.	N/A	N/A
6764019	12/2003	Kayahara et al.	N/A	N/A
6782385	12/2003	Natsumeda et al.	N/A	N/A
6813532	12/2003	Eryurek et al.	N/A	N/A
6816811	12/2003	Seem	N/A	N/A
6823680	12/2003	Jayanth	N/A	N/A
6826454	12/2003	Sulfstede	N/A	N/A
6865511	12/2004	Frerichs et al.	N/A	N/A
6925338	12/2004	Eryurek et al.	N/A	N/A
6986138	12/2005	Sakaguchi et al.	N/A	N/A
7031880	12/2005	Seem et al.	N/A	N/A
7401057	12/2007	Eder	N/A	N/A
7552467	12/2008	Lindsay	N/A	N/A
7627544	12/2008	Chkodrov et al.	N/A	N/A

7818249	12/2009	Lovejoy et al.	N/A	N/A
7889051	12/2010	Billig et al.	N/A	N/A
7996488	12/2010	Casabella et al.	N/A	N/A
8078330	12/2010	Brickfield et al.	N/A	N/A
8104044	12/2011	Scofield et al.	N/A	N/A
8229470	12/2011	Ranjan et al.	N/A	N/A
8401991	12/2012	Wu et al.	N/A	N/A
8495745	12/2012	Schrecker et al.	N/A	N/A
8516016	12/2012	Park et al.	N/A	N/A
8532808	12/2012	Drees et al.	N/A	N/A
8532839	12/2012	Drees et al.	N/A	N/A
8600556	12/2012	Nesler et al.	N/A	N/A
8635182	12/2013	Mackay	N/A	N/A
8682921	12/2013	Park et al.	N/A	N/A
8731724	12/2013	Drees et al.	N/A	N/A
8737334	12/2013	Ahn et al.	N/A	N/A
8738334	12/2013	Jiang et al.	N/A	N/A
8751487	12/2013	Byrne et al.	N/A	N/A
8788097	12/2013	Drees et al.	N/A	N/A
8805995	12/2013	Oliver	N/A	N/A
8843238	12/2013	Wenzel et al.	N/A	N/A
8874071	12/2013	Sherman et al.	N/A	N/A
8941465	12/2014	Pineau et al.	N/A	N/A
8990127	12/2014	Taylor	N/A	N/A
9070113	12/2014	Shafiee et al.	N/A	N/A
9116978	12/2014	Park et al.	N/A	N/A
9185095	12/2014	Moritz et al.	N/A	N/A
9189527	12/2014	Park et al.	N/A	N/A
9196009	12/2014	Drees et al.	N/A	N/A
9229966	12/2015	Aymeloglu et al.	N/A	N/A
9286582	12/2015	Drees et al.	N/A	N/A
9311807	12/2015	Schultz et al.	N/A	N/A
9344751	12/2015	Ream et al.	N/A	N/A
9354968	12/2015	Wenzel et al.	N/A	N/A
9507686	12/2015	Horn et al.	N/A	N/A
9524594	12/2015	Ouyang et al.	N/A	N/A
9558196	12/2016	Johnston et al.	N/A	N/A
9652813	12/2016	Gifford et al.	N/A	N/A
9753455	12/2016	Drees	N/A	N/A
9811249	12/2016	Chen et al.	N/A	N/A
9838844	12/2016	Emeis et al.	N/A	N/A
9886478	12/2017	Mukherjee	N/A	N/A
9948359	12/2017	Horton	N/A	N/A
10055114	12/2017	Shah et al.	N/A	N/A
10055206	12/2017	Park et al.	N/A	N/A
10116461	12/2017	Fairweather et al.	N/A	N/A
10169454	12/2018	Ait-Mokhtar et al.	N/A	N/A
10171297	12/2018	Stewart et al.	N/A	N/A
10171586	12/2018	Shaashua et al.	N/A	N/A
10187258	12/2018	Nagesh et al.	N/A	N/A
10514963	12/2018	Shrivastava et al.	N/A	N/A
10515098	12/2018	Park et al.	N/A	N/A
10534326	12/2019	Sridharan et al.	N/A	N/A
10536295	12/2019	Fairweather et al.	N/A	N/A

10564993	12/2019	Deutsch et al.	N/A	N/A
10705492	12/2019	Harvey	N/A	N/A
10708078	12/2019	Harvey	N/A	N/A
10760815	12/2019	Janakiraman et al.	N/A	N/A
10762475	12/2019	Song et al.	N/A	N/A
10824120	12/2019	Ahmed	N/A	N/A
10845771	12/2019	Harvey	N/A	N/A
10854194	12/2019	Park et al.	N/A	N/A
10862928	12/2019	Badawy et al.	N/A	N/A
10921760	12/2020	Harvey	N/A	N/A
10921972	12/2020	Park et al.	N/A	N/A
10969133	12/2020	Harvey	N/A	N/A
10986121	12/2020	Stockdale et al.	N/A	N/A
11016998	12/2020	Park et al.	N/A	N/A
11024292	12/2020	Park et al.	N/A	N/A
11038709	12/2020	Park et al.	N/A	N/A
11041650	12/2020	Li et al.	N/A	N/A
11054796	12/2020	Holaso	N/A	N/A
11070390	12/2020	Park et al.	N/A	N/A
11073976	12/2020	Park et al.	N/A	N/A
11108587	12/2020	Park et al.	N/A	N/A
11113295	12/2020	Park et al.	N/A	N/A
11229138	12/2021	Harvey et al.	N/A	N/A
11314726	12/2021	Park et al.	N/A	N/A
11314788	12/2021	Park et al.	N/A	N/A
11463322	12/2021	Sha et al.	N/A	N/A
11556105	12/2022	Cooley et al.	N/A	N/A
11561522	12/2022	Cooley et al.	N/A	N/A
11561523	12/2022	Cooley et al.	N/A	N/A
11573551	12/2022	Cooley et al.	N/A	N/A
11586167	12/2022	Cooley et al.	N/A	N/A
2002/0010562	12/2001	Schleiss et al.	N/A	N/A
2002/0016639	12/2001	Smith et al.	N/A	N/A
2002/0059229	12/2001	Natsumeda et al.	N/A	N/A
2002/0123864	12/2001	Eryurek et al.	N/A	N/A
2002/0147506	12/2001	Eryurek et al.	N/A	N/A
2002/0177909	12/2001	Fu et al.	N/A	N/A
2003/0005486	12/2002	Ridolfo et al.	N/A	N/A
2003/0014130	12/2002	Grumelart	N/A	N/A
2003/0073432	12/2002	Meade, II	N/A	N/A
2003/0158704	12/2002	Triginai et al.	N/A	N/A
2003/0171851	12/2002	Brickfield et al.	N/A	N/A
2003/0200059	12/2002	Ignatowski et al.	N/A	N/A
2004/0068390	12/2003	Saunders	N/A	N/A
2004/0128314	12/2003	Katibah et al.	N/A	N/A
2004/0133314	12/2003	Ehlers et al.	N/A	N/A
2004/0199360	12/2003	Friman et al.	N/A	N/A
2005/0055308	12/2004	Meyer et al.	N/A	N/A
2005/0108262	12/2004	Fawcett et al.	N/A	N/A
2005/0154494	12/2004	Ahmed	N/A	N/A
2005/0278703	12/2004	Lo et al.	N/A	N/A
2005/0283337	12/2004	Sayal	N/A	N/A
2006/0095521	12/2005	Patinkin	N/A	N/A
2006/0140207	12/2005	Eschbach et al.	N/A	N/A

2006/0184479	12/2005	Levine	N/A	N/A
2006/0200476	12/2005	Gottumukkala et al.	N/A	N/A
2006/0265751	12/2005	Cosquer et al.	N/A	N/A
2006/0271589	12/2005	Horowitz et al.	N/A	N/A
2007/0028179	12/2006	Levin et al.	N/A	N/A
2007/0203693	12/2006	Estes	N/A	N/A
2007/0261062	12/2006	Bansal et al.	N/A	N/A
2007/0273497	12/2006	Kuroda et al.	N/A	N/A
2007/0273610	12/2006	Baillot	N/A	N/A
2008/0034425	12/2007	Overcash et al.	N/A	N/A
2008/0094230	12/2007	Mock et al.	N/A	N/A
2008/0097816	12/2007	Freire et al.	N/A	N/A
2008/0186160	12/2007	Kim et al.	N/A	N/A
2008/0249756	12/2007	Chaisuparasmikul	N/A	N/A
2008/0252723	12/2007	Park	N/A	N/A
2008/0281472	12/2007	Podgorny et al.	N/A	N/A
2009/0195349	12/2008	Frader-Thompson et al.	N/A	N/A
2010/0045439	12/2009	Tak et al.	N/A	N/A
2010/0058248	12/2009	Park	N/A	N/A
2010/0131533	12/2009	Ortiz	N/A	N/A
2010/0274366	12/2009	Fata et al.	N/A	N/A
2010/0281387	12/2009	Holland et al.	N/A	N/A
2010/0286937	12/2009	Hedley et al.	N/A	N/A
2010/0324962	12/2009	Nesler et al.	N/A	N/A
2011/0015802	12/2010	Imes	N/A	N/A
2011/0047418	12/2010	Drees et al.	N/A	N/A
2011/0061015	12/2010	Drees et al.	N/A	N/A
2011/0071685	12/2010	Huneycutt et al.	N/A	N/A
2011/0077950	12/2010	Hughston	N/A	N/A
2011/0087650	12/2010	Mackay et al.	N/A	N/A
2011/0087988	12/2010	Ray et al.	N/A	N/A
2011/0088000	12/2010	Mackay	N/A	N/A
2011/0125737	12/2010	Pothering et al.	N/A	N/A
2011/0137853	12/2010	Mackay	N/A	N/A
2011/0153603	12/2010	Adiba et al.	N/A	N/A
2011/0154363	12/2010	Karmarkar	N/A	N/A
2011/0157357	12/2010	Weisensale et al.	N/A	N/A
2011/0178977	12/2010	Drees	N/A	N/A
2011/0191343	12/2010	Heaton et al.	N/A	N/A
2011/0205022	12/2010	Cavallaro et al.	N/A	N/A
2011/0218777	12/2010	Chen et al.	N/A	N/A
2012/0011126	12/2011	Park et al.	N/A	N/A
2012/0011141	12/2011	Park et al.	N/A	N/A
2012/0022698	12/2011	Mackay	N/A	N/A
2012/0062577	12/2011	Nixon	N/A	N/A
2012/0064923	12/2011	Imes et al.	N/A	N/A
2012/0083930	12/2011	Ilic et al.	N/A	N/A
2012/0100825	12/2011	Sherman et al.	N/A	N/A
2012/0101637	12/2011	Imes et al.	N/A	N/A
2012/0135759	12/2011	Imes et al.	N/A	N/A
2012/0136485	12/2011	Weber et al.	N/A	N/A
2012/0158633	12/2011	Eder	N/A	N/A
2012/0239797	12/2011	Agrawal et al.	N/A	N/A
2012/0259583	12/2011	Noboa et al.	N/A	N/A

2012/0272228	12/2011	Marndi et al.	N/A	N/A
2012/0278051	12/2011	Jiang et al.	N/A	N/A
2013/0007063	12/2012	Kalra et al.	N/A	N/A
2013/0038430	12/2012	Blower et al.	N/A	N/A
2013/0038707	12/2012	Cunningham et al.	N/A	N/A
2013/0060820	12/2012	Bulusu et al.	N/A	N/A
2013/0086497	12/2012	Ambuhl et al.	N/A	N/A
2013/0097706	12/2012	Titonis et al.	N/A	N/A
2013/0103221	12/2012	Raman et al.	N/A	N/A
2013/0167035	12/2012	Imes et al.	N/A	N/A
2013/0170710	12/2012	Kuoch et al.	N/A	N/A
2013/0204836	12/2012	Choi et al.	N/A	N/A
2013/0246916	12/2012	Reimann et al.	N/A	N/A
2013/0247205	12/2012	Schrecker et al.	N/A	N/A
2013/0262035	12/2012	Mills	N/A	N/A
2013/0275174	12/2012	Bennett et al.	N/A	N/A
2013/0275908	12/2012	Reichard	N/A	N/A
2013/0297050	12/2012	Reichard et al.	N/A	N/A
2013/0298244	12/2012	Kumar et al.	N/A	N/A
2013/0331995	12/2012	Rosen	N/A	N/A
2013/0338970	12/2012	Reggetti	N/A	N/A
2014/0032506	12/2013	Hoey et al.	N/A	N/A
2014/0059483	12/2013	Mairs et al.	N/A	N/A
2014/0081652	12/2013	Klindworth	N/A	N/A
2014/0135952	12/2013	Maehara	N/A	N/A
2014/0152651	12/2013	Chen et al.	N/A	N/A
2014/0172184	12/2013	Schmidt et al.	N/A	N/A
2014/0189861	12/2013	Gupta et al.	N/A	N/A
2014/0207282	12/2013	Angle et al.	N/A	N/A
2014/0258052	12/2013	Khuti et al.	N/A	N/A
2014/0269614	12/2013	Maguire et al.	N/A	N/A
2014/0277765	12/2013	Karimi et al.	N/A	N/A
2014/0278461	12/2013	Artz	N/A	N/A
2014/0327555	12/2013	Sager et al.	N/A	N/A
2015/0019174	12/2014	Kiff et al.	N/A	N/A
2015/0042240	12/2014	Aggarwal et al.	N/A	N/A
2015/0105917	12/2014	Sasaki et al.	N/A	N/A
2015/0145468	12/2014	Ma et al.	N/A	N/A
2015/0156031	12/2014	Fadell et al.	N/A	N/A
2015/0168931	12/2014	Jin	N/A	N/A
2015/0172300	12/2014	Cochenour	N/A	N/A
2015/0178421	12/2014	Borrelli et al.	N/A	N/A
2015/0185261	12/2014	Frader-Thompson et al.	N/A	N/A
2015/0186777	12/2014	Lecue et al.	N/A	N/A
2015/0202962	12/2014	Habashima et al.	N/A	N/A
2015/0204563	12/2014	Imes et al.	N/A	N/A
2015/0215389	12/2014	Spencer	707/741	H04L 67/1001
2015/0235267	12/2014	Steube et al.	N/A	N/A
2015/0241895	12/2014	Lu et al.	N/A	N/A
2015/0244730	12/2014	Vu et al.	N/A	N/A
2015/0244732	12/2014	Golshan et al.	N/A	N/A
2015/0261863	12/2014	Dey et al.	N/A	N/A
2015/0263900	12/2014	Polyakov et al.	N/A	N/A
2015/0286969	12/2014	Warner et al.	N/A	N/A

2015/0295796	12/2014	Hsiao et al.	N/A	N/A
2015/0304193	12/2014	Ishii et al.	N/A	N/A
2015/0316918	12/2014	Schleiss et al.	N/A	N/A
2015/0324422	12/2014	Elder	N/A	N/A
2015/0341212	12/2014	Hsiao et al.	N/A	N/A
2015/0348417	12/2014	Ignaczak et al.	N/A	N/A
2015/0379080	12/2014	Jochimski	N/A	N/A
2016/0011753	12/2015	McFarland et al.	N/A	N/A
2016/0033946	12/2015	Zhu et al.	N/A	N/A
2016/0035246	12/2015	Curtis	N/A	N/A
2016/0065601	12/2015	Gong et al.	N/A	N/A
2016/0070736	12/2015	Swan et al.	N/A	N/A
2016/0078229	12/2015	Gong et al.	N/A	N/A
2016/0090839	12/2015	Stolarczyk	N/A	N/A
2016/0119434	12/2015	Dong et al.	N/A	N/A
2016/0127712	12/2015	Alfredsson et al.	N/A	N/A
2016/0139752	12/2015	Shim et al.	N/A	N/A
2016/0163186	12/2015	Davidson et al.	N/A	N/A
2016/0170390	12/2015	Xie et al.	N/A	N/A
2016/0171862	12/2015	Das et al.	N/A	N/A
2016/0173816	12/2015	Huenerfauth et al.	N/A	N/A
2016/0179315	12/2015	Sarao et al.	N/A	N/A
2016/0179342	12/2015	Sarao et al.	N/A	N/A
2016/0179990	12/2015	Sarao et al.	N/A	N/A
2016/0195856	12/2015	Spero	N/A	N/A
2016/0212165	12/2015	Singla et al.	N/A	N/A
2016/0239660	12/2015	Azvine et al.	N/A	N/A
2016/0239756	12/2015	Aggour et al.	N/A	N/A
2016/0247129	12/2015	Song et al.	N/A	N/A
2016/0260063	12/2015	Harris et al.	N/A	N/A
2016/0313751	12/2015	Risbeck et al.	N/A	N/A
2016/0313752	12/2015	Przybylski	N/A	N/A
2016/0313902	12/2015	Hill et al.	N/A	N/A
2016/0350364	12/2015	Anicic et al.	N/A	N/A
2016/0357521	12/2015	Zhang et al.	N/A	N/A
2016/0357828	12/2015	Tobin et al.	N/A	N/A
2016/0358432	12/2015	Branscomb et al.	N/A	N/A
2016/0363336	12/2015	Roth et al.	N/A	N/A
2016/0370258	12/2015	Perez	N/A	N/A
2016/0378306	12/2015	Kresl et al.	N/A	N/A
2016/0379326	12/2015	Chan-Gove et al.	N/A	N/A
2017/0006135	12/2016	Siebel	N/A	N/A
2017/0011318	12/2016	Vigano et al.	N/A	N/A
2017/0017221	12/2016	Lamparter et al.	N/A	N/A
2017/0039255	12/2016	Raj et al.	N/A	N/A
2017/0052536	12/2016	Warner et al.	N/A	N/A
2017/0053441	12/2016	Nadumane et al.	N/A	N/A
2017/0063894	12/2016	Muddu et al.	N/A	N/A
2017/0068409	12/2016	Nair	N/A	N/A
2017/0070775	12/2016	Taxier et al.	N/A	N/A
2017/0075984	12/2016	Deshpande et al.	N/A	N/A
2017/0084168	12/2016	Janchookiat	N/A	N/A
2017/0090437	12/2016	Veeramani et al.	N/A	N/A
2017/0093700	12/2016	Gilley et al.	N/A	N/A

2017/0098086	12/2016	Hoernecke et al.	N/A	N/A
2017/0103327	12/2016	Penilla et al.	N/A	N/A
2017/0103403	12/2016	Chu et al.	N/A	N/A
2017/0123389	12/2016	Baez et al.	N/A	N/A
2017/0134415	12/2016	Muddu et al.	N/A	N/A
2017/0177715	12/2016	Chang et al.	N/A	N/A
2017/0180147	12/2016	Brandman et al.	N/A	N/A
2017/0188216	12/2016	Koskas et al.	N/A	N/A
2017/0212482	12/2016	Boettcher et al.	N/A	N/A
2017/0212668	12/2016	Shah et al.	N/A	N/A
2017/0220641	12/2016	Chi et al.	N/A	N/A
2017/0230930	12/2016	Frey	N/A	N/A
2017/0235817	12/2016	Deodhar et al.	N/A	N/A
2017/0251182	12/2016	Siminoff et al.	N/A	N/A
2017/0270124	12/2016	Nagano et al.	N/A	N/A
2017/0277769	12/2016	Pasupathy et al.	N/A	N/A
2017/0278003	12/2016	Liu	N/A	N/A
2017/0294132	12/2016	Colmenares	N/A	N/A
2017/0315522	12/2016	Kwon et al.	N/A	N/A
2017/0315697	12/2016	Jacobson et al.	N/A	N/A
2017/0322534	12/2016	Sinha et al.	N/A	N/A
2017/0323389	12/2016	Vavrasek	N/A	N/A
2017/0329289	12/2016	Kohn et al.	N/A	N/A
2017/0336770	12/2016	MacMillan	N/A	N/A
2017/0345287	12/2016	Fuller et al.	N/A	N/A
2017/0351957	12/2016	Lecue et al.	N/A	N/A
2017/0357225	12/2016	Asp et al.	N/A	N/A
2017/0357490	12/2016	Park et al.	N/A	N/A
2017/0357908	12/2016	Cabadi et al.	N/A	N/A
2018/0012159	12/2017	Kozloski et al.	N/A	N/A
2018/0013579	12/2017	Fairweather et al.	N/A	N/A
2018/0024520	12/2017	Sinha et al.	N/A	N/A
2018/0039238	12/2017	Gärtner et al.	N/A	N/A
2018/0048485	12/2017	Pelton et al.	N/A	N/A
2018/0069932	12/2017	Tiwari et al.	N/A	N/A
2018/0114140	12/2017	Chen et al.	N/A	N/A
2018/0137288	12/2017	Polyakov	N/A	N/A
2018/0157930	12/2017	Rutschman et al.	N/A	N/A
2018/0162400	12/2017	Abdar	N/A	N/A
2018/0176241	12/2017	Manadhata et al.	N/A	N/A
2018/0198627	12/2017	Mullins	N/A	N/A
2018/0202675	12/2017	Park et al.	N/A	N/A
2018/0203961	12/2017	Aisu et al.	N/A	N/A
2018/0239982	12/2017	Rutschman et al.	N/A	N/A
2018/0275625	12/2017	Park et al.	N/A	N/A
2018/0276962	12/2017	Butler et al.	N/A	N/A
2018/0292797	12/2017	Lamparter et al.	N/A	N/A
2018/0336785	12/2017	Ghannam et al.	N/A	N/A
2018/0356775	12/2017	Harvey	N/A	N/A
2018/0359111	12/2017	Harvey	N/A	N/A
2018/0364654	12/2017	Locke et al.	N/A	N/A
2019/0005025	12/2018	Malabarba	N/A	N/A
2019/0013023	12/2018	Pourmohammad et al.	N/A	N/A
2019/0025771	12/2018	Park et al.	N/A	N/A

2019/0026359	12/2018	Park et al.	N/A	N/A
2019/0037135	12/2018	Hedge	N/A	N/A
2019/0042988	12/2018	Brown et al.	N/A	N/A
2019/0088106	12/2018	Grundstrom	N/A	N/A
2019/0094824	12/2018	Xie et al.	N/A	N/A
2019/0096217	12/2018	Pourmohammad et al.	N/A	N/A
2019/0098113	12/2018	Park et al.	N/A	N/A
2019/0102840	12/2018	Perl et al.	N/A	N/A
2019/0121801	12/2018	Jethwa et al.	N/A	N/A
2019/0138512	12/2018	Pourmohammad et al.	N/A	N/A
2019/0147883	12/2018	Mellenthin et al.	N/A	N/A
2019/0158309	12/2018	Park et al.	N/A	N/A
2019/0163152	12/2018	Worrall et al.	N/A	N/A
2019/0226708	12/2018	Craft	N/A	N/A
2019/0258747	12/2018	Milev	N/A	N/A
2019/0268178	12/2018	Fairweather et al.	N/A	N/A
2019/0294978	12/2018	Sachs	N/A	N/A
2019/0310979	12/2018	Masuzaki et al.	N/A	N/A
2019/0354922	12/2018	Berti	N/A	G06K 7/1417
2019/0361412	12/2018	Park et al.	N/A	N/A
2019/0377306	12/2018	Harvey	N/A	N/A
2020/0050163	12/2019	Ludwig et al.	N/A	N/A
2020/0125043	12/2019	Park	N/A	N/A
2020/0133470	12/2019	Fala et al.	N/A	N/A
2020/0133472	12/2019	Fala et al.	N/A	N/A
2020/0133978	12/2019	Ramamurti et al.	N/A	N/A
2020/0142365	12/2019	Sharma et al.	N/A	N/A
2020/0226156	12/2019	Borra et al.	N/A	N/A
2020/0234220	12/2019	Ma	N/A	G08G 1/148
2020/0285203	12/2019	Thakur et al.	N/A	N/A
2020/0285657	12/2019	Ravizza et al.	N/A	N/A
2020/0336328	12/2019	Harvey	N/A	N/A
2020/0348632	12/2019	Harvey	N/A	N/A
2020/0387576	12/2019	Brett et al.	N/A	N/A
2020/0396208	12/2019	Brett et al.	N/A	N/A
2021/0042299	12/2020	Migliori	N/A	N/A
2021/0043221	12/2020	Yelchuru et al.	N/A	N/A
2021/0208546	12/2020	Locke et al.	N/A	N/A
2021/0311718	12/2020	Sinha et al.	N/A	N/A
2021/0325070	12/2020	Endel et al.	N/A	N/A
2021/0342961	12/2020	Winter et al.	N/A	N/A
2021/0373509	12/2020	Borah et al.	N/A	N/A
2021/0373510	12/2020	Borah et al.	N/A	N/A
2021/0381711	12/2020	Harvey et al.	N/A	N/A
2021/0381712	12/2020	Harvey et al.	N/A	N/A
2021/0382445	12/2020	Harvey et al.	N/A	N/A
2021/0383041	12/2020	Harvey et al.	N/A	N/A
2021/0383042	12/2020	Harvey et al.	N/A	N/A
2021/0383200	12/2020	Harvey et al.	N/A	N/A
2021/0383219	12/2020	Harvey et al.	N/A	N/A
2021/0383235	12/2020	Harvey et al.	N/A	N/A
2021/0383236	12/2020	Harvey et al.	N/A	N/A
2022/0004445	12/2021	George et al.	N/A	N/A
2022/0066402	12/2021	Harvey et al.	N/A	N/A

2022/0066405	12/2021	Harvey	N/A	N/A
2022/0066432	12/2021	Harvey et al.	N/A	N/A
2022/0066434	12/2021	Harvey et al.	N/A	N/A
2022/0066528	12/2021	Harvey et al.	N/A	N/A
2022/0066722	12/2021	Harvey et al.	N/A	N/A
2022/0066754	12/2021	Harvey et al.	N/A	N/A
2022/0066761	12/2021	Harvey et al.	N/A	N/A
2022/0067226	12/2021	Harvey et al.	N/A	N/A
2022/0067227	12/2021	Harvey et al.	N/A	N/A
2022/0067230	12/2021	Harvey et al.	N/A	N/A
2022/0069863	12/2021	Harvey et al.	N/A	N/A
2022/0070293	12/2021	Harvey et al.	N/A	N/A
2022/0121965	12/2021	Chatterji et al.	N/A	N/A
2022/0138684	12/2021	Harvey	N/A	N/A
2022/0147000	12/2021	Cooley et al.	N/A	N/A
2022/0150124	12/2021	Cooley et al.	N/A	N/A
2022/0215264	12/2021	Harvey et al.	N/A	N/A
2022/0397888	12/2021	Kamath	N/A	G05B 17/02
2022/0405327	12/2021	Swanson et al.	N/A	N/A
2022/0405668	12/2021	Swanson et al.	N/A	N/A
2023/0010757	12/2022	Preciado	N/A	N/A
2023/0071312	12/2022	Preciado et al.	N/A	N/A
2023/0076011	12/2022	Preciado et al.	N/A	N/A
2023/0083703	12/2022	Meiners	N/A	N/A
2023/0214555	12/2022	Harvey et al.	N/A	N/A
2023/0252205	12/2022	Harvey et al.	N/A	N/A
2023/0334200	12/2022	Meiners	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2019226217	12/2019	AU	N/A
2019226264	12/2019	AU	N/A
101415011	12/2008	CN	N/A
102136099	12/2010	CN	N/A
102136100	12/2010	CN	N/A
102650876	12/2011	CN	N/A
104040583	12/2013	CN	N/A
104603832	12/2014	CN	N/A
104919484	12/2014	CN	N/A
106204392	12/2015	CN	N/A
106406806	12/2016	CN	N/A
106960269	12/2016	CN	N/A
107147639	12/2016	CN	N/A
107598928	12/2017	CN	N/A
2 528 033	12/2011	EP	N/A
3 268 821	12/2017	EP	N/A
3 324 306	12/2017	EP	N/A
4 226 263	12/2022	EP	N/A
H10-049552	12/1997	JP	N/A
2003-162573	12/2002	JP	N/A
2007-018322	12/2006	JP	N/A
4073946	12/2007	JP	N/A
2008-107930	12/2007	JP	N/A
2013-152618	12/2012	JP	N/A

2014-044457	12/2013	JP	N/A
2016/0102923	12/2015	KR	N/A
WO-2009/020158	12/2008	WO	N/A
WO-2011/100255	12/2010	WO	N/A
WO-2013/050333	12/2012	WO	N/A
WO-2015/106702	12/2014	WO	N/A
WO-2015/145648	12/2014	WO	N/A
WO-2017/035536	12/2016	WO	N/A
WO-2017/192422	12/2016	WO	N/A
WO-2017/194244	12/2016	WO	N/A
WO-2017/205330	12/2016	WO	N/A
WO-2017/213918	12/2016	WO	N/A
WO-2018/132112	12/2017	WO	N/A
WO-2020/061621	12/2019	WO	N/A
WO-2022/042925	12/2021	WO	N/A
WO-2022/103812	12/2021	WO	N/A
WO-2022/103813	12/2021	WO	N/A
WO-2022/103820	12/2021	WO	N/A
WO-2022/103822	12/2021	WO	N/A
WO-2022/103824	12/2021	WO	N/A
WO-2022/103829	12/2021	WO	N/A
WO-2022/103831	12/2021	WO	N/A

OTHER PUBLICATIONS

U.S. Appl. No. 17/566,029, Passivelogic, Inc. cited by applicant

U.S. Appl. No. 17/567,275, Passivelogic, Inc. cited by applicant

U.S. Appl. No. 17/722,115, Passivelogic, Inc. cited by applicant

Balaji et al., “Brick: Metadata schema for portable smart building applications,” Applied Energy, 2018 (20 pages). cited by applicant

Balaji et al., “Brick: Metadata schema for portable smart building applications,” Applied Energy, Sep. 15, 2018, 3 pages, (Abstract). cited by applicant

Balaji et al., “Demo Abstract: Portable Queries Using the Brick Schema for Building Applications,” BuildSys '16, Palo Alto, CA, USA, Nov. 16-17, 2016 (2 pages). cited by applicant

Balaji, B. et al., “Brick: Towards a Unified Metadata Schema for Buildings.” BuildSys '16, Palo Alto, CA, USA, Nov. 16-17, 2016 (10 pages). cited by applicant

Bhattacharya et al., “Short Paper: Analyzing Metadata Schemas for Buildings—The Good, The Bad and The Ugly,” BuildSys '15, Seoul, South Korea, Nov. 4-5, 2015 (4 pages). cited by applicant

Bhattacharya, A., “Enabling Scalable Smart-Building Analytics,” Electrical Engineering and Computer Sciences, University of California at Berkeley, Technical Report No. UCB/EECS-2016-201, Dec. 15, 2016 (121 pages). cited by applicant

Brick, “Brick Schema: Building Blocks for Smart Buildings,” URL: chrome-extension://efaidnbmnnnibpcajpcglclefindmkaj/https://www.memoori.com/wp-content/uploads/2016/06/Brick_Schema_Whitepaper.pdf, Mar. 2019 (17 pages). cited by applicant

Brick, “Brick: Towards a Unified Metadata Schema for Buildings,” URL: chrome-extension://efaidnbmnnnibpcajpcglclefindmkaj/https://brickschema.org/papers/Brick_BuildSys_Presentation.pdf, Presented at BuildSys '16, Nov. 2016 (46 pages). cited by applicant

Brick, “Metadata Schema for Buildings,” URL: https://brickschema.org/docs/Brick-Leaflet.pdf, retrieved from internet Dec. 24, 2019 (3 pages). cited by applicant

Chinese Office Action on CN Appl. Ser. No. 201780003995.9 dated Apr. 8, 2021 (21 pages with English language translation). cited by applicant

Chinese Office action on CN Appl. Ser. No. 201780043400.2 dated Apr. 25, 2021 (15 pages with English language translation). cited by applicant

Curry, E. et al., “Linking building data in the cloud: Integrating cross-domain building data using linked data.” Advanced Engineering Informatics, 2013, 27 (pp. 206-219). cited by applicant

Digital Platform Litigation Documents Part 1, includes cover letter, dismissal of case DDE-1-21-cv-01796, IPR2023-00022 (documents filed Jan. 26, 2023-Oct. 7, 2022), and IPR2023-00085 (documents filed Jan. 26, 2023-Oct. 20, 2022) (748 pages total). cited by applicant

Digital Platform Litigation Documents Part 10, includes DDE-1-21-cv-01796 (documents filed Nov. 1, 2022-Dec. 22, 2021 (1795 pages total). cited by applicant

Digital Platform Litigation Documents Part 2, includes IPR2023-00085 (documents filed Oct. 20, 2022) (172 pages total). cited by applicant

Digital Platform Litigation Documents Part 3, includes IPR2023-00085 (documents filed Oct. 20, 2022) and IPR2023-00170 (documents filed Nov. 28, 2022-Nov. 7, 2022) (397 pages total). cited by applicant

Digital Platform Litigation Documents Part 4, includes IPR2023-00170 (documents filed Nov. 7, 2022) and IPR2023-00217 (documents filed Jan. 18, 2023-Nov. 15, 2022) (434 pages total). cited by applicant

Digital Platform Litigation Documents Part 5, includes IPR2023-00217 (documents filed Nov. 15, 2022) and IPR2023-00257 (documents filed Jan. 25, 2023-Nov. 23, 2022) (316 pages total). cited by applicant

Digital Platform Litigation Documents Part 6, includes IPR2023-00257 (documents filed Nov. 23, 2022) and IPR2023-00346 (documents filed Jan. 3, 2023-Dec. 13, 2022) (295 pages total). cited by applicant

Digital Platform Litigation Documents Part 7, includes IPR2023-00346 (documents filed Dec. 13, 2022) and IPR2023-00347 (documents filed Jan. 3, 2023-Dec. 13, 2022) (217 pages total). cited by applicant

Digital Platform Litigation Documents Part 8, includes IPR2023-00347 (documents filed Dec. 13, 2022), EDTX-2-22-cv-00243 (documents filed Sep. 20, 2022-Jun. 29, 2022), and DDE-1-21-cv-01796 (documents filed Feb. 3, 2023-Jan. 10, 2023 (480 pages total). cited by applicant

Digital Platform Litigation Documents Part 9, includes DDE-1-21-cv-01796 (documents filed Jan. 10, 2023-Nov.1, 2022 (203 pages total). cited by applicant

El Kaed, C. et al., “Building management insights driven by a multi-system semantic representation approach,” 2016 IEEE 3rd World Forum on Internet of Things (WF-IoT), Dec. 12-14, 2016, (pp. 520-525). cited by applicant

Ellis, C. et al., “Creating a room connectivity graph of a building from per-room sensor units.” BuildSys '12, Toronto, ON, Canada, Nov. 6, 2012 (7 pages). cited by applicant

Extended European Search Report on EP Application No. 18196948.6 dated Apr. 10, 2019 (9 pages). cited by applicant

Fierro et al., “Beyond a House of Sticks: Formalizing Metadata Tags with Brick,” BuildSys '19, New York, NY, USA, Nov. 13-14, 2019 (10 pages). cited by applicant

Fierro et al., “Dataset: An Open Dataset and Collection Tool for BMS Point Labels,” Data'19, New York, NY, USA, Nov. 10, 2019 (3 pages). cited by applicant

Fierro et al., “Design and Analysis of a Query Processor for Brick,” ACM Transactions on Sensor Networks, Jan. 2018, vol. 1, No. 1, (25 pages). cited by applicant

Fierro et al., “Design and Analysis of a Query Processor for Brick,” BuildSys '17, Delft, Netherlands, Nov. 8-9, 2017 (10 pages). cited by applicant

Fierro et al., “Mortar: An Open Testbed for Portable Building Analytics,” BuildSys '18, Shenzhen, China, Nov. 7-8, 2018 (10 pages). cited by applicant

Fierro et al., “Why Brick is a Game Changer for Smart Buildings,” URL: https://brickschema.org/papers/Brick_Memoori_Webinar_Presentation.pdf, Memoori Webinar, 2019 (67 pages). cited by applicant

Fierro, “Writing Portable Building Analytics with the Brick Metadata Schema,” UC Berkeley, ACM E-Energy, 2019 (39 pages). cited by applicant

Fierro, G., “Design of an Effective Ontology and Query Processor Enabling Portable Building Applications,” Electrical Engineering and Computer Sciences, University of California at Berkeley, Technical Report No. UCB/EECS-2019-106, Jun. 27, 2019 (118 pages). cited by applicant

File History for U.S. Appl. No. 12/776,159, filed May 7, 2010 (722 pages). cited by applicant

Final Conference Program, ACM BuildSys 2016, Stanford, CA, USA, Nov. 15-17, 2016 (7 pages). cited by applicant

Gao et al., “A large-scale evaluation of automated metadata inference approaches on sensors from air handling units,” Advanced Engineering Informatics, 2018, 37 (pp. 14-30). cited by applicant

Harvey, T., “Quantum Part 3: The Tools of Autonomy, How PassiveLogic's Quantum Creator and Autonomy Studio software works,” URL:

<https://www.automatedbuildings.com/news/jan22/articles/passive/211224010000passive.html>, Jan. 2022 (7 pages). cited by applicant

Harvey, T., “Quantum: The Digital Twin Standard for Buildings,” URL: <https://www.automatedbuildings.com/news/feb21/articles/passivelogic/210127124501passivelogic.html>, Feb. 2021 (6 pages). cited by applicant

Hu, S. et al., “Building performance optimisation: A hybrid architecture for the integration of contextual information and time-series data,” *Automation in Construction*, 2016, 70 (pp. 51-61). cited by applicant

International Search Report and Written Opinion for PCT Appl. Ser. No. PCT/US2017/013831 dated Mar. 31, 2017 (14 pages). cited by applicant

International Search Report and Written Opinion for PCT Appl. Ser. No. PCT/US2017/035524 dated Jul. 24, 2017 (14 pages). cited by applicant

International Search Report and Written Opinion on PCT/US2017/052060, mailed Oct. 5, 2017, 11 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2017/052633, mailed Oct. 23, 2017, 9 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2017/052829, mailed Nov. 27, 2017, 24 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2018/024068, mailed Jun. 15, 2018, 22 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2018/052971, dated Mar. 1, 2019, 19 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2018/052974, mailed Dec. 19, 2018, 13 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2018/052975, mailed Jan. 2, 2019, 13 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2018/052994, mailed Jan. 7, 2019, 15 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2019/015481, dated May 17, 2019, 15 pages. cited by applicant

International Search Report and Written Opinion on PCT/US2020/058381, dated Jan. 27, 2021, 30 pages. cited by applicant

Japanese Office Action on JP Appl. Ser. No. 2018-534963 dated May 11, 2021 (16 pages with English language translation). cited by applicant

Koh et al., “Plaster: An Integration, Benchmark, and Development Framework for Metadata Normalization Methods,” *BuildSys '18*, Shenzhen, China, Nov. 7-8, 2018 (10 pages). cited by applicant

Koh et al., “Scrabble: Transferrable Semi-Automated Semantic Metadata Normalization using Intermediate Representation,” *BuildSys '18*, Shenzhen, China, Nov. 7-8, 2018 (10 pages). cited by applicant

Koh et al., “Who can Access What, and When?” *BuildSys '19*, New York, NY, USA, Nov. 13-14, 2019 (4 pages). cited by applicant

Li et al., “Event Stream Processing with Out-of-Order Data Arrival,” *International Conferences on Distributed Computing Systems*, 2007, (8 pages). cited by applicant

Nissin Electric Co., Ltd., “Smart power supply system (SPSS),” Outline of the scale verification plan, Nissin Electric Technical Report, Japan, Apr. 23, 2014, vol. 59, No. 1 (23 pages). cited by applicant

Passivelogic, “Explorer: Digital Twin Standard for Autonomous Systems. Made interactive.” URL: <https://passivelogic.com/software/quantum-explorer/>, retrieved from internet Jan. 4, 2023 (13 pages). cited by applicant

Passivelogic, “Quantum: The Digital Twin Standard for Autonomous Systems, A physics-based ontology for next-generation control and AI.” URL: <https://passivelogic.com/software/quantum-standard/>, retrieved from internet Jan. 4, 2023 (20 pages). cited by applicant

Quantum Alliance, “Quantum Explorer Walkthrough,” 2022, (7 pages) (screenshots from video). cited by applicant

Results of the Partial International Search for PCT/US2018/052971, dated Jan. 3, 2019, 3 pages. cited by applicant

Sinha, Sudhi and Al Huraimel, Khaled, “Reimagining Businesses with AI” John Wiley & Sons, Inc., Hoboken,

NJ, USA, First ed. published 2020 (156 pages). cited by applicant

Sinha, Sudhi R. and Park, Youngchoon, "Building an Effective IoT Ecosystem for Your Business," Johnson Controls International, Springer International Publishing, 2017 (286 pages). cited by applicant

Sinha, Sudhi, "Making Big Data Work for Your Business: A guide to effective Big Data analytics," Impackt Publishing LTD., Birmingham, UK, Oct. 2014 (170 pages). cited by applicant

The Virtual Nuclear Tourist, "Calvert Cliffs Nuclear Power Plant," URL: <http://www.nucleartourist.com/us/calvert.htm>, Jan. 11, 2006 (2 pages). cited by applicant

University of California at Berkeley, EECS Department, "Enabling Scalable Smart-Building Analytics," URL: <https://www2.eecs.berkeley.edu/Pubs/TechRpts/2016/EECS-2016-201.html>, retrieved from internet Feb. 15, 2023 (7 pages). cited by applicant

Van Hoof, Bert, "Announcing Azure Digital Twins: Create digital replicas of spaces and infrastructure using cloud, AI and IoT," URL: <https://azure.microsoft.com/en-us/blog/announcing-azure-digital-twins-create-digital-replicas-of-spaces-and-infrastructure-using-cloud-ai-and-iot/>, Sep. 24, 2018 (11 pages). cited by applicant

W3C, "SPARQL: Query Language for RDF," located on The Wayback Machine, URL: <https://web.archive.org/web/20161230061728/http://www.w3.org/TR/rdf-sparql-query/>), retrieved from internet Nov. 15, 2022 (89 pages). cited by applicant

Wei et al., "Development and Implementation of Software Gateways of Fire Fighting Subsystem Running on EBI," Control, Automation and Systems Engineering, IITA International Conference on, IEEE, Jul. 2009 (pp. 9-12). cited by applicant

White et al., "Reduce building maintenance costs with AWS IoT TwinMaker Knowledge Graph," The Internet of Things on AWS—Official Blog, URL: <https://aws.amazon.com/blogs/iot/reduce-building-maintenance-costs-with-aws-iot-twinmaker-knowledge-graph/>, Nov. 18, 2022 (10 pages). cited by applicant

Zhou, Q. et al., "Knowledge-infused and Consistent Complex Event Processing over Real-time and Persistent Streams," Further Generation Computer Systems, 2017, 76 (pp. 391-406). cited by applicant

European Search Report on EP Appl. No. 23165820.4 dated Aug. 22, 2023 (11 pages). cited by applicant

Gao et al., "An Anomaly Detection Framework for Digital Twin Driven Cyber-Physical Systems", ICCPS '21, Nashville, TN, USA, May 19-21, 2021 (pp. 44-54). cited by applicant

Office Action on European Appl. No. EP 23165820.4 dated Mar. 19, 2025 (6 pages). cited by applicant

Wang et al., "Linking energy-cyber-physical systems with occupancy prediction and interpretation through WiFi probe-based ensemble classification," Applied Energy, 2019, vol. 236 (pp. 55-69). cited by applicant

Yao et al., "State of the art review on model predictive control (MPC) in Heating Ventilation and Air-conditioning (HVAC) field," Building and Environment, 2021, vol. 200 (18 pages). cited by applicant

Primary Examiner: Gofman; Alex

Attorney, Agent or Firm: Foley & Lardner LLP

Background/Summary

CROSS-REFERENCE TO RELATED PATENT APPLICATION (1) This application is a continuation of U.S. patent application Ser. No. 17/534,821 filed Nov. 24, 2021, the entirety of which is incorporated by reference herein.

BACKGROUND

(1) The present disclosure relates generally to the management of building systems and devices of a building. The present disclosure relates more particularly to managing information of building systems and controlling the building systems.

(2) A building can include various types of building subsystems, e.g., heating, ventilation, and air conditioning (HVAC) systems, security systems, fire response systems, access control systems, etc. In some cases, a central cloud system may manage the information of the various systems and control the various systems of the building. However, in some cases, local computing resources may be able to provide better real-time management of systems of a building. However, it may be difficult to manage the storage of information and operational responsibilities across local computing resources and cloud

computing resources.

SUMMARY

(3) One implementation of the present disclosure is a method including receiving, by one or more processing circuits, at least a portion of building data describing entities of a building and relationships between the entities, the entities represent at least one of a device, a space, the building, a point, or a person. The method includes generating, by the one or more processing circuits, a first digital twin based on the building data, wherein a first system stores the first digital twin and a second system stores a second digital twin generated based on the building data. The first digital twin includes first entities of the entities and first relationships of the relationships between the first entities, wherein the second digital twin includes second entities of the entities and second relationships of the relationships between the second entities, wherein the first digital twin includes a relationship that forms a connection between the first digital twin and the second digital twin by linking a first entity of the first entities of the first digital twin and a second entity of the second entities of the second digital twin. The method includes performing, by the one or more processing circuits, one or more operations based on at least one of the first digital twin, the second digital twin, or the relationship that forms the connection between the first digital twin and the second digital twin.

(4) In some embodiments, the first system and the second system are at least one of a cloud system stored off premises from the building or an edge system stored on a premises of the building.

(5) In some embodiments, the one or more processing circuits include at least one of a first processing circuit of the first system or a second processing circuit of the second system.

(6) In some embodiments, the one or more processing circuits include a processing circuit of the first system. In some embodiments, performing the one or more operations includes executing, by the processing circuit of the first system, an agent stored by the first system, the agent configured to perform the one or more operations based on the first digital twin and ingesting, by the agent, a result of the one or more operations into the first digital twin.

(7) In some embodiments, the second digital twin includes a second relationship that forms another connection between the second digital twin and the first digital twin by linking the second entity of the second entities of the second digital twin and the first entity of the first entities of the first digital twin.

(8) In some embodiments, the one or more processing circuits include a first processing circuit of the first system and a second processing circuit of the second system. In some embodiments, the method includes performing an onboarding to generate the first digital twin and the second digital twin by receiving, by the first processing circuit of the first system, first building data of the building data, generating, by the first processing circuit of the first system, the first digital twin based on the first building data, receiving, by the second processing circuit of the second system, second building data of the building data, generating, by the second processing circuit of the second system, the second digital twin based on the second data, identifying, by the first processing circuit of the first system, the relationship that forms the connection between the first digital twin and the second digital twin by communicating with the second system, and causing, by the first processing circuit of the first system, the first digital twin to include the relationship.

(9) In some embodiments, the first digital twin includes a first graph structure including first nodes representing the first entities and first edges between the first nodes representing the first relationships between the first entities. In some embodiments, the second digital twin includes a second graph structure including second nodes representing the second entities and second edges between the second nodes representing the second relationships between the second entities.

(10) In some embodiments, the first graph structure includes a node identifying the second graph structure. In some embodiments, the first graph structure includes an edge between a first node of the first nodes representing the first entity and the node identifying the second graph structure, wherein the edge represents the relationship between the first entity of the first entities of the first digital twin and the second entity of the second entities of the second digital twin.

(11) In some embodiments, the edge includes data including a first indication of the first entity, a second indication that the first entity is stored by the first graph structure, a third indication of the second entity, a fourth indication that the second entity is stored by the second graph structure, and a fifth indication of a relationship type of relationship types describing the relationship between the first entity of the first entities of the first digital twin and the second entity of the second entities of the second digital twin.

(12) In some embodiments, the one or more processing circuits include a processing circuit of the first

system. In some embodiments, performing the one or more operations includes querying, by the processing circuit of the first system, the first digital twin based on a query, identifying, by the processing circuit of the first system, that a query result of the query is stored in the second digital twin based on the relationship that forms the connection between the first digital twin and the second digital twin, communicating, by the processing circuit of the first system, the query to the second system, receiving, by the processing circuit of the first system, the query result, wherein the query result is based on the second system querying the second digital twin based on the query, and performing, by the processing circuit of the first system, the one or more operations based on the query result.

(13) In some embodiments, the one or more processing circuits include a second processing circuit of the second system. In some embodiments, the method includes receiving, by the second processing circuit of the second system, the query from the processing circuit of the first system, querying, by the second processing circuit of the second system, the second digital twin with the query to generate the query result, and communicating, by the second processing circuit of the first system, the query result to the processing circuit of the first system.

(14) In some embodiments, the method includes identifying, by the processing circuit of the first system, that the second system stores the second digital twin based on a lookup table. In some embodiments, the lookup table stores indications of digital twins including the first digital twin and the second digital twin and links between the digital twins and systems that store the digital twins, the systems including the first system and the second system.

(15) Another implementation of the present disclosure is a system including one or more memory devices storing instructions thereon that, when executed by one or more processors, cause the one or more processors to receive at least a portion of building data describing entities of a building and relationships between the entities, the entities represent at least one of a device, a space, the building, a point, or a person. The instructions cause the one or more processors to generate a first digital twin based on the building data, wherein the system stores the first digital twin and a second system stores a second digital twin generated based on the building data, wherein the first digital twin includes first entities of the entities and first relationships of the relationships between the first entities, wherein the second digital twin includes second entities of the entities and second relationships of the relationships between the second entities, wherein the first digital twin includes a relationship that forms a connection between the first digital twin and the second digital twin by linking a first entity of the first entities of the first digital twin and a second entity of the second entities of the second digital twin. The instructions cause the one or more processors to perform one or more operations based on at least one of the first digital twin, the second digital twin, or the relationship that forms the connection between the first digital twin and the second digital twin.

(16) In some embodiments, the system and the second system are at least one of a cloud system stored off premises from the building or an edge system stored on a premises of the building.

(17) In some embodiments, the instructions cause the one or more processors to execute an agent stored by the first system, the agent configured to perform the one or more operations based on the first digital twin and ingest a result of the one or more operations into the first digital twin.

(18) In some embodiments, the second digital twin includes a second relationship that forms another connection between the second digital twin and the first digital twin by linking the second entity of the second entities of the second digital twin and the first entity of the first entities of the first digital twin.

(19) In some embodiments, the first digital twin includes a first graph structure including first nodes representing the first entities and first edges between the first nodes representing the first relationships between the first entities. In some embodiments, the second digital twin includes a second graph structure including second nodes representing the second entities and second edges between the second nodes representing the second relationships between the second entities.

(20) In some embodiments, the instructions cause the one or more processors to query the first digital twin based on a query, identify that a query result of the query is stored in the second digital twin based on the relationship that forms the connection between the first digital twin and the second digital twin, communicate the query to the second system; receive the query result, wherein the query result is based on the second system querying the second digital twin based on the query, and perform the one or more operations based on the query result.

(21) Another implementation of the present disclosure is one or more computer readable medium storing

instructions thereon that, when executed by one or more processors, cause the one or more processors to receive at least a portion of building data describing entities of a building and relationships between the entities, the entities represent at least one of a device, a space, the building, a point, or a person. The instructions cause the one or more processors to generate a first digital twin based on the building data, wherein a first system stores the first digital twin and a second system stores a second digital twin generated based on the building data, wherein the first digital twin includes first entities of the entities and first relationships of the relationships between the first entities, wherein the second digital twin includes second entities of the entities and second relationships of the relationships between the second entities, wherein the first digital twin includes a relationship that forms a connection between the first digital twin and the second digital twin by linking a first entity of the first entities of the first digital twin and a second entity of the second entities of the second digital twin. The instructions cause the one or more processors to perform one or more operations based on at least one of the first digital twin, the second digital twin, or the relationship that forms the connection between the first digital twin and the second digital twin.

(22) In some embodiments, the first digital twin includes a first graph structure including first nodes representing the first entities and first edges between the first nodes representing the first relationships between the first entities. In some embodiments, the second digital twin includes a second graph structure including second nodes representing the second entities and second edges between the second nodes representing the second relationships between the second entities.

(23) In some embodiments, the instructions cause the first system to perform one or more particular operations based on data received from the second system, the data based on the second digital twin, identify a communication issue between the first system and the second system, and perform the one or more particular operations based on second data determined based on the first digital twin.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Various objects, aspects, features, and advantages of the disclosure will become more apparent and better understood by referring to the detailed description taken in conjunction with the accompanying drawings, in which like reference characters identify corresponding elements throughout. In the drawings, like reference numbers generally indicate identical, functionally similar, and/or structurally similar elements.
- (2) FIG. 1 is a block diagram of a building data platform including an edge platform, a cloud platform, and a twin manager, according to an exemplary embodiment.
- (3) FIG. 2 is a graph projection of the twin manager of FIG. 1 including application programming interface (API) data, capability data, policy data, and services, according to an exemplary embodiment.
- (4) FIG. 3 is another graph projection of the twin manager of FIG. 1 including application programming interface (API) data, capability data, policy data, and services, according to an exemplary embodiment.
- (5) FIG. 4 is a graph projection of the twin manager of FIG. 1 including equipment and capability data for the equipment, according to an exemplary embodiment.
- (6) FIG. 5 is a block diagram of a system for managing a digital twin where an artificial intelligence agent can be executed to infer information for an entity of a graph, according to an exemplary embodiment.
- (7) FIG. 6 is a flow diagram of a process for executing an artificial intelligence agent to infer and/or predict information, according to an exemplary embodiment.
- (8) FIG. 7 is a flow diagram an agent of a digital twin executing a trigger rule and an action rule, according to an exemplary embodiment.
- (9) FIG. 8 is a block diagram of a system where a clean air optimization (CAO) AI service and an energy prediction model (EPM) AI service operate to make inferences and predictions based on data of a knowledge graph, according to an exemplary embodiment.
- (10) FIG. 9 is a block diagram of a CAO AI agent and an EPM AI agent operating against the knowledge graph of FIG. 26 to generate inferences and predictions, according to an exemplary embodiment.
- (11) FIG. 10 is a block diagram of the edge platform of FIG. 1 performing event enrichment at the edge before the events are communicated to the cloud, according to an exemplary embodiment.
- (12) FIG. 11 is a block diagram a building graph divided and distributed across a cloud system and an edge

system, according to an exemplary embodiment.

(13) FIG. 12 is a block diagram of the twin manager of FIG. 1 and an edge device storing the building graph of FIG. 11 divided and distributed across the twin manager and the edge device, according to an exemplary embodiment.

(14) FIG. 13 is a block diagram of the distributed building graph generated across the twin manager and a building controller during an onboarding process, according to an exemplary embodiment.

(15) FIG. 14 is a flow diagram of a process of querying a building graph divided into multiple building graphs distributed across cloud and edge systems, according to an exemplary embodiment.

(16) FIG. 15 is a flow diagram of a process of querying a building graph divided into multiple building graphs distributed across cloud and edge systems by an agent, and ingesting information back into the building graph by the agent, according to an exemplary embodiment.

(17) FIG. 16 is a flow diagram of a process of querying a building graph divided into multiple graphs distributed across a cloud system and an edge platform and using information queried from the building graph to perform one or more operations, according to an exemplary embodiment.

(18) FIG. 17 is a flow diagram of a process where the distributed building graph is generated across the twin manager and a building controller during onboarding, according to an exemplary embodiment.

DETAILED DESCRIPTION

Overview

(19) Referring generally to the FIGURES, systems and methods for distributing a digital twin of a building across cloud and edge devices are shown, according to various exemplary embodiments. A digital twin can be a virtual representation of a building and/or an entity of the building (e.g., space, piece of equipment, occupant, etc.). Furthermore, the digital twin can represent a service performed in a building, e.g., facility management, clean air optimization, energy prediction, equipment maintenance, etc.

(20) In some embodiments, the digital twin can include an information data store and a connector. The information data store can store the information describing the entity that the digital twin operates for (e.g., attributes of the entity, measurements associated with the entity, control points or commands of the entity, etc.). In some embodiments, the data store can be a graph including various nodes and edges. The connector can be a software component that provides telemetry from the entity (e.g., physical device) to the information store. Furthermore, the digital twin can include an artificial intelligence (AI), e.g., an AI agent. The AI can be one or more rules based engines, machine learning algorithms, and/or models that operate based on information of the information data store and output information. The AI agent can run against the digital twin. In some embodiments, the AI agent can run against a common data model, e.g., a BRICK model, and can be easily implemented in various different buildings, e.g., against various different building models of the system format.

(21) In some embodiments, the AI agent for the digital twin can call an AI service to determine inferences and/or predict future data values. In some embodiments, the predictions are potential future states. In some embodiments, the predictions predict a timeseries of a data point into the future. The predictions could be predicted indoor temperature for an hour, inferred future air quality from 15 minute air quality readings, etc. In some embodiments, the digital twin can store predicted and/or inferred information in a graph data store as a node in the graph data store related to an entity that the digital twin represents or otherwise operates for. In some embodiments, the digital twin, or other digital twins, can operate against the predicted and/or inferred data, e.g., operate to construct and implement control algorithms for operating equipment of a building based on predicted future data points of the building.

(22) Some systems may implement a digital twin on a cloud platform only. However, such a system may not be able to handle a hybrid implementation of a digital twin where the digital twin is stored across a one or more cloud and/or edge systems. In some embodiments, the building system described herein can distribute the digital twin across one or more cloud and/or edge systems. In some embodiments, operational capabilities of a digital twin can be distributed and/or implemented by the building system across the one or more cloud and/or edge systems.

(23) In some embodiments, a building graph (which may be part of a digital twin or make up the digital twin) can be distributed across the one or more cloud and/or edge systems. The building graph can be divided into one or more interrelated building graphs that are distributed. In this regard, information pertaining to a particular edge system can be stored in a building graph specific to the edge system and interrelated to other building graphs stored in other cloud and/or edge systems. In this regard, algorithms,

analytics, computing elements, etc. of the particular edge system can run against the local building graph.

(24) In some embodiments, each building graph distributed across the one or more edge and/or cloud systems can be connected to each other. For example, a first building graph may store one or more nodes or edges that connect the first building graph to a second building graph stored by a second system. In this regard, a first system can traverse the nodes or edges of the first building graph that the first system stores to identify a second building graph stored by a second system, the second building graph storing information needed to fill a query of the first system. The first system can send the query to the identified second system and receive a result of the query filled by the second system based on the second building graph.

(25) In some embodiments, the divided and distributed building graph can help separate operating concerns through localizing operations. In some embodiments, the building system can distribute responsibilities to various edge and/or cloud systems. Because each edge system handles its own responsibilities with its own graph, the edge system may not need to delay data management or control for passing data or states between the edge system and a central cloud platform. In some embodiments, the portion of the building graph stored by each edge and/or cloud system can include its own format extensions and/or features. For example, an edge system that manages video surveillance systems may store its graph in a formatted extended with specific data types, attributes, relationship types etc. specific to video surveillance systems while a second HVAC system may store its graph in the same format but with extensions for specific data types, attributes, relationship types, etc. specific to HVAC systems.

(26) In some embodiments, distributing and/or dividing the building graph across multiple cloud and/or edge systems can resolve security issues. In some embodiments, data may be more secure if the data stays within a building and does not need to be transferred via one or more external networks from the building to a cloud system. Because the edge systems can store their own building graphs and perform operations based on the edge systems, the data of the building may not need to be transferred outside the building and thus is less likely to be obtained by other entities.

(27) The divided building graph can, in some embodiments, provide flexibility in where digital twin computations are performed. For example, because each cloud and/or edge system can include the building graph that each cloud and/or edge system needs, each cloud and/or edge system can locally run computational functions against the building graph that the system stores. In some embodiments, functions such as agents can run locally within each system. Examples of agents and agent related processing can be found in U.S. patent application Ser. No. 17/354,436 filed Jun. 22, 2021, U.S. patent application Ser. No. 17/354,338 filed Jun. 22, 2021, U.S. patent application Ser. No. 17/148,851 filed Jan. 14, 2021, U.S. patent application Ser. No. 15/723,624 filed Oct. 3, 2017, U.S. patent application Ser. No. 16/008,885 filed Jun. 14, 2018, U.S. patent application Ser. No. 16/143,243 filed Sep. 26, 2018, U.S. patent application Ser. No. 16/533,499 filed Aug. 6, 2019, and U.S. patent application Ser. No. 16/142,859 filed Aug. 26, 2018, the entirety of which is incorporated by reference herein.

(28) Referring now to FIG. 1, a building data platform **100** including an edge platform **102**, a cloud platform **106**, and a twin manager **108** are shown, according to an exemplary embodiment. The edge platform **102**, the cloud platform **106**, and the twin manager **108** can each be separate services deployed on the same or different computing systems. In some embodiments, the cloud platform **106** and the twin manager **108** are implemented in off premises computing systems, e.g., outside a building. The edge platform **102** can be implemented on-premises, e.g., within the building. However, any combination of on-premises and off-premises components of the building data platform **100** can be implemented.

(29) The building data platform **100** includes applications **110**. The applications **110** can be various applications that operate to manage the building subsystems **122**. The applications **110** can be remote or on-premises applications (or a hybrid of both) that run on various computing systems. The applications **110** can include an alarm application **168** configured to manage alarms for the building subsystems **122**. The applications **110** include an assurance application **170** that implements assurance services for the building subsystems **122**. In some embodiments, the applications **110** include an energy application **172** configured to manage the energy usage of the building subsystems **122**. The applications **110** include a security application **174** configured to manage security systems of the building.

(30) In some embodiments, the applications **110** and/or the cloud platform **106** interacts with a user device **176**. In some embodiments, a component or an entire application of the applications **110** runs on the user device **176**. The user device **176** may be a laptop computer, a desktop computer, a smartphone, a tablet,

and/or any other device with an input interface (e.g., touch screen, mouse, keyboard, etc.) and an output interface (e.g., a speaker, a display, etc.).

(31) The applications **110**, the twin manager **108**, the cloud platform **106**, and the edge platform **102** can be implemented on one or more computing systems, e.g., on processors and/or memory devices. For example, the edge platform **102** includes processor(s) **118** and memories **120**, the cloud platform **106** includes processor(s) **124** and memories **126**, the applications **110** include processor(s) **164** and memories **166**, and the twin manager **108** includes processor(s) **148** and memories **150**.

(32) The processors can be a general purpose or specific purpose processors, an application specific integrated circuit (ASIC), one or more field programmable gate arrays (FPGAs), a group of processing components, or other suitable processing components. The processors may be configured to execute computer code and/or instructions stored in the memories or received from other computer readable media (e.g., CDROM, network storage, a remote server, etc.).

(33) The memories can include one or more devices (e.g., memory units, memory devices, storage devices, etc.) for storing data and/or computer code for completing and/or facilitating the various processes described in the present disclosure. The memories can include random access memory (RAM), read-only memory (ROM), hard drive storage, temporary storage, non-volatile memory, flash memory, optical memory, or any other suitable memory for storing software objects and/or computer instructions. The memories can include database components, object code components, script components, or any other type of information structure for supporting the various activities and information structures described in the present disclosure. The memories can be communicably connected to the processors and can include computer code for executing (e.g., by the processors) one or more processes described herein.

(34) The edge platform **102** can be configured to provide connection to the building subsystems **122**. The edge platform **102** can receive messages from the building subsystems **122** and/or deliver messages to the building subsystems **122**. The edge platform **102** includes one or multiple gateways, e.g., the gateways **112-116**. The gateways **112-116** can act as a gateway between the cloud platform **106** and the building subsystems **122**. The gateways **112-116** can be the gateways described in U.S. Provisional Patent Application No. 62/951,897 filed Dec. 20, 2019, the entirety of which is incorporated by reference herein. In some embodiments, the applications **110** can be deployed on the edge platform **102**. In this regard, lower latency in management of the building subsystems **122** can be realized.

(35) The edge platform **102** can be connected to the cloud platform **106** via a network **104**. The network **104** can communicatively couple the devices and systems of building data platform **100**. In some embodiments, the network **104** is at least one of and/or a combination of a Wi-Fi network, a wired Ethernet network, a ZigBee network, a Bluetooth network, and/or any other wireless network. The network **104** may be a local area network or a wide area network (e.g., the Internet, a building WAN, etc.) and may use a variety of communications protocols (e.g., BACnet, IP, LON, etc.). The network **104** may include routers, modems, servers, cell towers, satellites, and/or network switches. The network **104** may be a combination of wired and wireless networks.

(36) The cloud platform **106** can be configured to facilitate communication and routing of messages between the applications **110**, the twin manager **108**, the edge platform **102**, and/or any other system. The cloud platform **106** can include a platform manager **128**, a messaging manager **140**, a command processor **136**, and an enrichment manager **138**. In some embodiments, the cloud platform **106** can facilitate messaging between the building data platform **100** via the network **104**.

(37) The messaging manager **140** can be configured to operate as a transport service that controls communication with the building subsystems **122** and/or any other system, e.g., managing commands to devices (C2D), commands to connectors (C2C) for external systems, commands from the device to the cloud (D2C), and/or notifications. The messaging manager **140** can receive different types of data from the applications **110**, the twin manager **108**, and/or the edge platform **102**. The messaging manager **140** can receive change on value data **142**, e.g., data that indicates that a value of a point has changed. The messaging manager **140** can receive timeseries data **144**, e.g., a time correlated series of data entries each associated with a particular time stamp. Furthermore, the messaging manager **140** can receive command data **146**. All of the messages handled by the cloud platform **106** can be handled as an event, e.g., the data **142-146** can each be packaged as an event with a data value occurring at a particular time (e.g., a temperature measurement made at a particular time).

(38) The cloud platform **106** includes a command processor **136**. The command processor **136** can be

configured to receive commands to perform an action from the applications **110**, the building subsystems **122**, the user device **176**, etc. The command processor **136** can manage the commands, determine whether the commanding system is authorized to perform the particular commands, and communicate the commands to the commanded system, e.g., the building subsystems **122** and/or the applications **110**. The commands could be a command to change an operational setting that control environmental conditions of a building, a command to run analytics, etc.

(39) The cloud platform **106** includes an enrichment manager **138**. The enrichment manager **138** can be configured to enrich the events received by the messaging manager **140**. The enrichment manager **138** can be configured to add contextual information to the events. The enrichment manager **138** can communicate with the twin manager **108** to retrieve the contextual information. In some embodiments, the contextual information is an indication of information related to the event. For example, if the event is a timeseries temperature measurement of a thermostat, contextual information such as the location of the thermostat (e.g., what room), the equipment controlled by the thermostat (e.g., what VAV), etc. can be added to the event. In this regard, when a consuming application, e.g., one of the applications **110** receives the event, the consuming application can operate based on the data of the event, the temperature measurement, and also the contextual information of the event.

(40) The enrichment manager **138** can solve a problem that when a device produces a significant amount of information, the information may contain simple data without context. An example might include the data generated when a user scans a badge at a badge scanner of the building subsystems **122**. This physical event can generate an output event including such information as “DeviceBadgeScannerID,” “BadgeID,” and/or “Date/Time.” However, if a system sends this data to a consuming application, e.g., Consumer A and a Consumer B, each customer may need to call the building data platform knowledge service to query information with queries such as, “What space, build, floor is that badge scanner in?” or “What user is associated with that badge?”

(41) By performing enrichment on the data feed, a system can be able to perform inferences on the data. A result of the enrichment may be transformation of the message “DeviceBadgeScannerId, BadgeId, Date/Time,” to “Region, Building, Floor, Asset, DeviceId, BadgeId, UserName, EmployeeId, Date/Time Scanned.” This can be a significant optimization, as a system can reduce the number of calls by $1/n$, where n is the number of consumers of this data feed.

(42) By using this enrichment, a system can also have the ability to filter out undesired events. If there are 100 building in a campus that receive 100,000 events per building each hour, but only 1 building is actually commissioned, only $1/10$ of the events are enriched. By looking at what events are enriched and what events are not enriched, a system can do traffic shaping of forwarding of these events to reduce the cost of forwarding events that no consuming application wants or reads.

(43) An example of an event received by the enrichment manager **138** may be:

(44) TABLE-US-00001 { “id”: “someguid”, “eventType”: “Device_Heartbeat”, “eventTime”: “2018-01-27T00:00:00+00:00” “eventValue”: 1, “deviceId”: “someguid” }

(45) An example of an enriched event generated by the enrichment manager **138** may be:

(46) TABLE-US-00002 { “id”: “someguid”, “eventType”: “Device_Heartbeat”, “eventTime”: “2018-01-27T00:00:00+00:00” “eventValue”: 1, “deviceId”: “someguid”, “buildingName”: “Building-48”, “buildingID”: “SomeGuid”, “panelID”: “SomeGuid”, “panelName”: “Building-48-Panel-13”, “cityID”: 371, “cityName”: “Milwaukee”, “stateID”: 48, “stateName”: “Wisconsin (WI)”, “countryID”: 1, “countryName”: “United States” }

(47) By receiving enriched events, an application of the applications **110** can be able to populate and/or filter what events are associated with what areas. Furthermore, user interface generating applications can generate user interfaces that include the contextual information based on the enriched events.

(48) The cloud platform **106** includes a platform manager **128**. The platform manager **128** can be configured to manage the users and/or subscriptions of the cloud platform **106**. For example, what subscribing building, user, and/or tenant utilizes the cloud platform **106**. The platform manager **128** includes a provisioning service **130** configured to provision the cloud platform **106**, the edge platform **102**, and the twin manager **108**. The platform manager **128** includes a subscription service **132** configured to manage a subscription of the building, user, and/or tenant while the entitlement service **134** can track entitlements of the buildings, users, and/or tenants.

(49) The twin manager **108** can be configured to manage and maintain a digital twin. The digital twin can

be a digital representation of the physical environment, e.g., a building. The twin manager **108** can include a change feed generator **152**, a schema and ontology **154**, a graph projection manager **156**, a policy manager **158**, an entity, relationship, and event database **160**, and a graph projection database **162**.

(50) The graph projection manager **156** can be configured to construct graph projections and store the graph projections in the graph projection database **162**. Examples of graph projections are shown in FIGS. **11-13**. Entities, relationships, and events can be stored in the database **160**. The graph projection manager **156** can retrieve entities, relationships, and/or events from the database **160** and construct a graph projection based on the retrieved entities, relationships and/or events. In some embodiments, the database **160** includes an entity-relationship collection for multiple subscriptions.

(51) In some embodiment, the graph projection manager **156** generates a graph projection for a particular user, application, subscription, and/or system. In this regard, the graph projection can be generated based on policies for the particular user, application, and/or system in addition to an ontology specific for that user, application, and/or system. In this regard, an entity could request a graph projection and the graph projection manager **156** can be configured to generate the graph projection for the entity based on policies and an ontology specific to the entity. The policies can indicate what entities, relationships, and/or events the entity has access to. The ontology can indicate what types of relationships between entities the requesting entity expects to see, e.g., floors within a building, devices within a floor, etc. Another requesting entity may have an ontology to see devices within a building and applications for the devices within the graph.

(52) The graph projections generated by the graph projection manager **156** and stored in the graph projection database **162** can be a knowledge graph and is an integration point. For example, the graph projections can represent floor plans and systems associated with each floor. Furthermore, the graph projections can include events, e.g., telemetry data of the building subsystems **122**. The graph projections can show application services as nodes and API calls between the services as edges in the graph. The graph projections can illustrate the capabilities of spaces, users, and/or devices. The graph projections can include indications of the building subsystems **122**, e.g., thermostats, cameras, VAVs, etc. The graph projection database **162** can store graph projections that keep up a current state of a building.

(53) The graph projections of the graph projection database **162** can be digital twins of a building. Digital twins can be digital replicas of physical entities that enable an in-depth analysis of data of the physical entities and provide the potential to monitor systems to mitigate risks, manage issues, and utilize simulations to test future solutions. Digital twins can play an important role in helping technicians find the root cause of issues and solve problems faster, in supporting safety and security protocols, and in supporting building managers in more efficient use of energy and other facilities resources. Digital twins can be used to enable and unify security systems, employee experience, facilities management, sustainability, etc.

(54) In some embodiments the enrichment manager **138** can use a graph projection of the graph projection database **162** to enrich events. In some embodiments, the enrichment manager **138** can identify nodes and relationships that are associated with, and are pertinent to, the device that generated the event. For example, the enrichment manager **138** could identify a thermostat generating a temperature measurement event within the graph. The enrichment manager **138** can identify relationships between the thermostat and spaces, e.g., a zone that the thermostat is located in. The enrichment manager **138** can add an indication of the zone to the event.

(55) Furthermore, the command processor **136** can be configured to utilize the graph projections to command the building subsystems **122**. The command processor **136** can identify a policy for a commanding entity within the graph projection to determine whether the commanding entity has the ability to make the command. For example, the command processor **136**, before allowing a user to make a command, determine, based on the graph projection database **162**, to determine that the user has a policy to be able to make the command.

(56) In some embodiments, the policies can be conditional based policies. For example, the building data platform **100** can apply one or more conditional rules to determine whether a particular system has the ability to perform an action. In some embodiments, the rules analyze a behavioral based biometric. For example, a behavioral based biometric can indicate normal behavior and/or normal behavior rules for a system. In some embodiments, when the building data platform **100** determines, based on the one or more conditional rules, that an action requested by a system does not match a normal behavior, the building data

platform **100** can deny the system the ability to perform the action and/or request approval from a higher level system.

(57) For example, a behavior rule could indicate that a user has access to log into a system with a particular IP address between 8 A.M. through 5 P.M. However, if the user logs in to the system at 7 P.M., the building data platform **100** may contact an administrator to determine whether to give the user permission to log in.

(58) The change feed generator **152** can be configured to generate a feed of events that indicate changes to the digital twin, e.g., to the graph. The change feed generator **152** can track changes to the entities, relationships, and/or events of the graph. For example, the change feed generator **152** can detect an addition, deletion, and/or modification of a node or edge of the graph, e.g., changing the entities, relationships, and/or events within the database **160**. In response to detecting a change to the graph, the change feed generator **152** can generate an event summarizing the change. The event can indicate what nodes and/or edges have changed and how the nodes and edges have changed. The events can be posted to a topic by the change feed generator **152**.

(59) The change feed generator **152** can implement a change feed of a knowledge graph. The building data platform **100** can implement a subscription to changes in the knowledge graph. When the change feed generator **152** posts events in the change feed, subscribing systems or applications can receive the change feed event. By generating a record of all changes that have happened, a system can stage data in different ways, and then replay the data back in whatever order the system wishes. This can include running the changes sequentially one by one and/or by jumping from one major change to the next. For example, to generate a graph at a particular time, all change feed events up to the particular time can be used to construct the graph.

(60) The change feed can track the changes in each node in the graph and the relationships related to them, in some embodiments. If a user wants to subscribe to these changes and the user has proper access, the user can simply submit a web API call to have sequential notifications of each change that happens in the graph. A user and/or system can replay the changes one by one to reinstitute the graph at any given time slice. Even though the messages are “thin” and only include notification of change and the reference “id/seq id,” the change feed can keep a copy of every state of each node and/or relationship so that a user and/or system can retrieve those past states at any time for each node. Furthermore, a consumer of the change feed could also create dynamic “views” allowing different “snapshots” in time of what the graph looks like from a particular context. While the twin manager **108** may contain the history and the current state of the graph based upon schema evaluation, a consumer can retain a copy of that data, and thereby create dynamic views using the change feed.

(61) The schema and ontology **154** can define the message schema and graph ontology of the twin manager **108**. The message schema can define what format messages received by the messaging manager **140** should have, e.g., what parameters, what formats, etc. The ontology can define graph projections, e.g., the ontology that a user wishes to view. For example, various systems, applications, and/or users can be associated with a graph ontology. Accordingly, when the graph projection manager **156** generates an graph projection for a user, system, or subscription, the graph projection manager **156** can generate a graph projection according to the ontology specific to the user. For example, the ontology can define what types of entities are related in what order in a graph, for example, for the ontology for a subscription of “Customer A,” the graph projection manager **156** can create relationships for a graph projection based on the rule: Regioncustom characterBuildingcustom characterFloorcustom characterSpacecustom characterAsset

(62) For the ontology of a subscription of “Customer B,” the graph projection manager **156** can create relationships based on the rule: Buildingcustom characterFloorcustom characterAsset

(63) The policy manager **158** can be configured to respond to requests from other applications and/or systems for policies. The policy manager **158** can consult a graph projection to determine what permissions different applications, users, and/or devices have. The graph projection can indicate various permissions that different types of entities have and the policy manager **158** can search the graph projection to identify the permissions of a particular entity. The policy manager **158** can facilitate fine grain access control with user permissions. The policy manager **158** can apply permissions across a graph, e.g., if “user can view all data associated with floor **1**” then they see all subsystem data for that floor, e.g., surveillance cameras, HVAC devices, fire detection and response devices, etc.

(64) The twin manager **108** includes a query manager **165** and a twin function manager **167**. The query manager **164** can be configured to handle queries received from a requesting system, e.g., the user device **176**, the applications **110**, and/or any other system. The query manager **165** can receive queries that include query parameters and context. The query manager **165** can query the graph projection database **162** with the query parameters to retrieve a result. The query manager **165** can then cause an event processor, e.g., a twin function, to operate based on the result and the context. In some embodiments, the query manager **165** can select the twin function based on the context and/or perform operates based on the context.

(65) The twin function manager **167** can be configured to manage the execution of twin functions. The twin function manager **167** can receive an indication of a context query that identifies a particular data element and/or pattern in the graph projection database **162**. Responsive to the particular data element and/or pattern occurring in the graph projection database **162** (e.g., based on a new data event added to the graph projection database **162** and/or change to nodes or edges of the graph projection database **162**, the twin function manager **167** can cause a particular twin function to execute. The twin function can execute based on an event, context, and/or rules. The event can be data that the twin function executes against. The context can be information that provides a contextual description of the data, e.g., what device the event is associated with, what control point should be updated based on the event, etc.

(66) Referring now to FIG. 2, a graph projection **200** of the twin manager **108** including application programming interface (API) data, capability data, policy data, and services is shown, according to an exemplary embodiment. The graph projection **200** includes nodes **202-240** and edges **250-272**. The nodes **202-240** and the edges **250-272** are defined according to the key **201**. The nodes **202-240** represent different types of entities, devices, locations, points, persons, policies, and software services (e.g., API services). The edges **250-272** represent relationships between the nodes **202-240**, e.g., dependent calls, API calls, inferred relationships, and schema relationships (e.g., BRICK relationships).

(67) The graph projection **200** includes a device hub **202** which may represent a software service that facilitates the communication of data and commands between the cloud platform **106** and a device of the building subsystems **122**, e.g., door actuator **214**. The device hub **202** is related to a connector **204**, an external system **206**, and a digital asset “Door Actuator” **208** by edge **250**, edge **252**, and edge **254**.

(68) The cloud platform **106** can be configured to identify the device hub **202**, the connector **204**, the external system **206** related to the door actuator **214** by searching the graph projection **200** and identifying the edges **250-254** and edge **258**. The graph projection **200** includes a digital representation of the “Door Actuator,” node **208**. The digital asset “Door Actuator” **208** includes a “DeviceNameSpace” represented by node **207** and related to the digital asset “Door Actuator” **208** by the “Property of Object” edge **256**.

(69) The “Door Actuator” **214** has points and timeseries. The “Door Actuator” **214** is related to “Point A” **216** by a “has_a” edge **260**. The “Door Actuator” **214** is related to “Point B” **218** by a “has_a” edge **258**. Furthermore, timeseries associated with the points A and B are represented by nodes “TS” **220** and “TS” **222**. The timeseries are related to the points A and B by “has_a” edge **264** and “has_a” edge **262**. The timeseries “TS” **220** has particular samples, sample **210** and **212** each related to “TS” **220** with edges **268** and **266** respectively. Each sample includes a time and a value. Each sample may be an event received from the door actuator that the cloud platform **106** ingests into the entity, relationship, and event database **160**, e.g., ingests into the graph projection **200**.

(70) The graph projection **200** includes a building **234** representing a physical building. The building includes a floor represented by floor **232** related to the building **234** by the “has a” edge from the building **234** to the floor **232**. The floor has a space indicated by the edge “has_a” **270** between the floor **232** and the space **230**. The space has particular capabilities, e.g., is a room that can be booked for a meeting, conference, private study time, etc. Furthermore, the booking can be canceled. The capabilities for the floor **232** are represented by capabilities **228** related to space **230** by edge **280**. The capabilities **228** are related to two different commands, command “book room” **224** and command “cancel booking” **226** related to capabilities **228** by edge **284** and edge **282** respectively.

(71) If the cloud platform **106** receives a command to book the space represented by the node, space **230**, the cloud platform **106** can search the graph projection **200** for the capabilities for the **228** related to the space **230** to determine whether the cloud platform **106** can book the room.

(72) In some embodiments, the cloud platform **106** could receive a request to book a room in a particular building, e.g., the building **234**. The cloud platform **106** could search the graph projection **200** to identify

spaces that have the capabilities to be booked, e.g., identify the space **230** based on the capabilities **228** related to the space **230**. The cloud platform **106** can reply to the request with an indication of the space and allow the requesting entity to book the space **230**.

(73) The graph projection **200** includes a policy **236** for the floor **232**. The policy **236** is related set for the floor **232** based on a “To Floor” edge **274** between the policy **236** and the floor **232**. The policy **236** is related to different roles for the floor **232**, read events **238** via edge **276** and send command **240** via edge **278**. The policy **236** is set for the entity **203** based on has edge **251** between the entity **203** and the policy **236**.

(74) The twin manager **108** can identify policies for particular entities, e.g., users, software applications, systems, devices, etc. based on the policy **236**. For example, if the cloud platform **106** receives a command to book the space **230**. The cloud platform **106** can communicate with the twin manager **108** to verify that the entity requesting to book the space **230** has a policy to book the space. The twin manager **108** can identify the entity requesting to book the space as the entity **203** by searching the graph projection **200**. Furthermore, the twin manager **108** can further identify the edge has **251** between the entity **203** and the policy **236** and the edge **1178** between the policy **236** and the command **240**.

(75) Furthermore, the twin manager **108** can identify that the entity **203** has the ability to command the space **230** based on the edge **1174** between the policy **236** and the edge **270** between the floor **232** and the space **230**. In response to identifying the entity **203** has the ability to book the space **230**, the twin manager **108** can provide an indication to the cloud platform **106**.

(76) Furthermore, if the entity makes a request to read events for the space **230**, e.g., the sample **210** and the sample **212**, the twin manager **108** can identify the edge has **251** between the entity **203** and the policy **236**, the edge **1178** between the policy **236** and the read events **238**, the edge **1174** between the policy **236** and the floor **232**, the “has_a” edge **270** between the floor **232** and the space **230**, the edge **268** between the space **230** and the door actuator **214**, the edge **260** between the door actuator **214** and the point A **216**, the “has_a” edge **264** between the point A **216** and the TS **220**, and the edges **268** and **266** between the TS **220** and the samples **210** and **212** respectively.

(77) Referring now to FIG. **3**, a graph projection **300** of the twin manager **108** including application programming interface (API) data, capability data, policy data, and services is shown, according to an exemplary embodiment. The graph projection **300** includes the nodes and edges described in the graph projection **200** of FIG. **2**. The graph projection **300** includes a connection broker **310** related to capabilities **228** by edge **398a**. The connection broker **310** can be a node representing a software application configured to facilitate a connection with another software application. In some embodiments, the cloud platform **106** can identify the system that implements the capabilities **228** by identifying the edge **398a** between the capabilities **228** and the connection broker **310**.

(78) The connection broker **310** is related to an agent that optimizes a space **356** via edge **398b**. The agent represented by the node **356** can book and cancel bookings for the space represented by the node **230** based on the edge **398b** between the connection broker **310** and the node **356** and the edge **398a** between the capabilities **228** and the connection broker **310**.

(79) The connection broker **310** is related to a cluster **308** by edge **398c**. Cluster **308** is related to connector B **302** via edge **398e** and connector A **306** via edge **398d**. The connector A **306** is related to an external subscription service **304**. A connection broker **310** is related to cluster **308** via an edge **311** representing a rest call that the connection broker represented by node **310** can make to the cluster represented by cluster **308**.

(80) The connection broker **310** is related to a virtual meeting platform **312** by an edge **354**. The node **312** represents an external system that represents a virtual meeting platform. The connection broker represented by node **310** can represent a software component that facilitates a connection between the cloud platform **106** and the virtual meeting platform represented by node **312**. When the cloud platform **106** needs to communicate with the virtual meeting platform represented by the node **312**, the cloud platform **106** can identify the edge **354** between the connection broker **310** and the virtual meeting platform **312** and select the connection broker represented by the node **310** to facilitate communication with the virtual meeting platform represented by the node **312**.

(81) A capabilities node **318** can be connected to the connection broker **310** via edge **360**. The capabilities **318** can be capabilities of the virtual meeting platform represented by the node **312** and can be related to the node **312** through the edge **360** to the connection broker **310** and the edge **354** between the connection

broker **310** and the node **312**. The capabilities **318** can define capabilities of the virtual meeting platform represented by the node **312**. The node **320** is related to capabilities **318** via edge **362**. The capabilities may be an invite bob command represented by node **316** and an email bob command represented by node **314**. The capabilities **318** can be linked to a node **320** representing a user, Bob. The cloud platform **106** can facilitate email commands to send emails to the user Bob via the email service represented by the node **304**. The node **304** is related to the connect a node **306** via edge **398f**. Furthermore, the cloud platform **106** can facilitate sending an invite for a virtual meeting via the virtual meeting platform represented by the node **312** linked to the node **318** via the edge **358**.

(82) The node **320** for the user Bob can be associated with the policy **236** via the “has” edge **364**. Furthermore, the node **320** can have a “check policy” edge **366** with a portal node **324**. The device API node **328** has a check policy edge **370** to the policy node **236**. The portal node **324** has an edge **368** to the policy node **236**. The portal node **324** has an edge **323** to a node **326** representing a user input manager (UIM). The portal node **324** is related to the UIM node **326** via an edge **323**. The UIM node **326** has an edge **323** to a device API node **328**. The UIM node **326** is related to the door actuator node **214** via edge **372**. The door actuator node **214** has an edge **374** to the device API node **328**. The door actuator **214** has an edge **335** to the connector virtual object **334**. The device hub **332** is related to the connector virtual object via edge **380**. The device API node **328** can be an API for the door actuator **214**. The connector virtual object **334** is related to the device API node **328** via the edge **331**.

(83) The device API node **328** is related to a transport connection broker **330** via an edge **329**. The transport connection broker **330** is related to a device hub **332** via an edge **378**. The device hub represented by node **332** can be a software component that hands the communication of data and commands for the door actuator **214**. The cloud platform **106** can identify where to store data within the graph projection **300** received from the door actuator by identifying the nodes and edges between the points **216** and **218** and the device hub node **332**. Similarly, the cluster **308** can identify commands for the door actuator that can be facilitated by the device hub represented by the node **332**, e.g., by identifying edges between the device hub node **332** and an open door node **352** and an lock door node **350**. The door actuator **114** has an edge “has mapped an asset” **280** between the node **214** and a capabilities node **348**. The capabilities node **348** and the nodes **352** and **350** are linked by edges **396** and **394**.

(84) The device hub **332** is linked to a cluster **336** via an edge **384**. The cluster **336** is linked to connector A **340** and connector B **338** by edges **386** and the edge **389**. The connector A **340** and the connector B **338** is linked to an external system **344** via edges **388** and **390**. The external system **344** is linked to a door actuator **342** via an edge **392**.

(85) Referring now to FIG. 4, a graph projection **400** of the twin manager **108** including equipment and capability data for the equipment is shown, according to an exemplary embodiment. The graph projection **400** includes nodes **402-456** and edges **360-498f**. The cloud platform **106** can search the graph projection **400** to identify capabilities of different pieces of equipment.

(86) A building node **404** represents a particular building that includes two floors. A floor 1 node **402** is linked to the building node **404** via edge **460** while a floor 2 node **406** is linked to the building node **404** via edge **462**. The floor 2 includes a particular room **2023** represented by edge **464** between floor 2 node **406** and room **2023** node **408**. Various pieces of equipment are included within the room **2023**. A light represented by light node **416**, a bedside lamp node **414**, a bedside lamp node **412**, and a hallway light node **410** are related to room **2023** node **408** via edge **466**, edge **472**, edge **470**, and edge **468**.

(87) The light represented by light node **416** is related to a light connector **426** via edge **484**. The light connector **426** is related to multiple commands for the light represented by the light node **416** via edges **484**, **486**, and **488**. The commands may be a brightness setpoint **424**, an on command **425**, and a hue setpoint **428**. The cloud platform **106** can receive a request to identify commands for the light represented by the light **416** and can identify the nodes **424-428** and provide an indication of the commands represented by the node **424-428** to the requesting entity. The requesting entity can then send commands for the commands represented by the nodes **424-428**.

(88) The bedside lamp node **414** is linked to a bedside lamp connector **481** via an edge **413**. The connector **481** is related to commands for the bedside lamp represented by the bedside lamp node **414** via edges **492**, **496**, and **494**. The command nodes are a brightness setpoint node **432**, an on command node **434**, and a color command **436**. The hallway light **410** is related to a hallway light connector **446** via an edge **498d**. The hallway light connector **446** is linked to multiple commands for the hallway light node **410** via edges

498g, **498f**, and **498e**. The commands are represented by an on command node **452**, a hue setpoint node **450**, and a light bulb activity node **448**.

(89) The graph projection **400** includes a name space node **422** related to a server A node **418** and a server B node **420** via edges **474** and **476**. The name space node **422** is related to the bedside lamp connector **481**, the bedside lamp connector **444**, and the hallway light connector **446** via edges **482**, **480**, and **478**. The bedside lamp connector **444** is related to commands, e.g., the color command node **440**, the hue setpoint command **438**, a brightness setpoint command **456**, and an on command **454** via edges **498c**, **498b**, **498a**, and **498**.

(90) Referring now to FIG. 5, a system **500** for managing a digital twin where an artificial intelligence agent can be executed to infer and/or predict information for an entity of a graph is shown, according to an exemplary embodiment. The system **500** can be components of the building data platform **100**, e.g., components run on the processors and memories of the edge platform **102**, the cloud platform **106**, the twin manager **108**, and/or the applications **110**. The system **500** can, in some implementations, implement a digital twin with artificial intelligence.

(91) A digital twin (or a shadow) may be a computing entity that describes a physical thing (e.g., a building, spaces of a building, devices of a building, people of the building, equipment of the building, etc.) through modeling the physical thing through a set of attributes that define the physical thing. A digital twin can refer to a digital replica of physical assets (a physical device twin) and can be extended to store processes, people, places, systems that can be used for various purposes. The digital twin can include both the ingestion of information and actions learned and executed through artificial intelligence agents.

(92) In FIG. 5, the digital twin can be a graph **529** managed by the twin manager **108** and/or artificial intelligence agents **570**. In some embodiments, the digital twin is the combination of the graph **529** with the artificial intelligence agents **570**. In some embodiments, the digital twin enables the creation of a chronological time-series database of telemetry events for analytical purposes. In some embodiments, the graph **529** uses the BRICK schema.

(93) The twin manager **108** stores the graph **529** which may be a graph data structure including various nodes and edges interrelating the nodes. The graph **529** may be the same as, or similar to, the graph projections described herein with reference to FIGS. 1-4. The graph **529** includes nodes **510-526** and edges **528-546**. The graph **529** includes a building node **526** representing a building that has a floor indicated by the “has” edge **546** to the floor node **522**. The floor node **522** is related to a zone node **510** via a “has” edge **544** indicating that the floor represented by the node **522** has a zone represented by the zone **510**.

(94) The floor node **522** is related to the zone node **518** by the “has” edge **540** indicating that the floor represented by the floor node **522** has another zone represented by the zone node **518**. The floor node **522** is related to another zone node **524** via a “has” edge **542** representing that the floor represented by the floor node **522** has a third zone represented by the zone node **524**.

(95) The graph **529** includes an AHU node **514** representing an AHU of the building represented by the building node **526**. The AHU node **514** is related by a “supplies” edge **530** to the VAV node **512** to represent that the AHU represented by the AHU node **514** supplies air to the VAV represented by the VAV node **512**. The AHU node **514** is related by a “supplies” edge **536** to the VAV node **520** to represent that the AHU represented by the AHU node **514** supplies air to the VAV represented by the VAV node **520**. The AHU node **514** is related by a “supplies” edge **532** to the VAV node **516** to represent that the AHU represented by the AHU node **514** supplies air to the VAV represented by the VAV node **516**.

(96) The VAV node **516** is related to the zone node **518** via the “serves” edge **534** to represent that the VAV represented by the VAV node **516** serves (e.g., heats or cools) the zone represented by the zone node **518**. The VAV node **520** is related to the zone node **524** via the “serves” edge **538** to represent that the VAV represented by the VAV node **520** serves (e.g., heats or cools) the zone represented by the zone node **524**. The VAV node **512** is related to the zone node **510** via the “serves” edge **528** to represent that the VAV represented by the VAV node **512** serves (e.g., heats or cools) the zone represented by the zone node **510**.

(97) Furthermore, the graph **529** includes an edge **533** related to a timeseries node **564**. The timeseries node **564** can be information stored within the graph **529** and/or can be information stored outside the graph **529** in a different database (e.g., a timeseries database). In some embodiments, the timeseries node **564** stores timeseries data (or any other type of data) for a data point of the VAV represented by the VAV node **516**. The data of the timeseries node **564** can be aggregated and/or collected telemetry data of the timeseries node **564**.

(98) Furthermore, the graph **529** includes an edge **537** related to a timeseries node **566**. The timeseries node **566** can be information stored within the graph **529** and/or can be information stored outside the graph **529** in a different database (e.g., a timeseries database). In some embodiments, the timeseries node **566** stores timeseries data (or any other type of data) for a data point of the VAV represented by the VAV node **516**. The data of the timeseries node **564** can be inferred information, e.g., data inferred by one of the artificial intelligence agents **570** and written into the timeseries node **564** by the artificial intelligence agent **570**. In some embodiments, the timeseries **564** and/or **566** are stored in the graph **529** but are stored as references to timeseries data stored in a timeseries database.

(99) The twin manager **108** includes various software components. For example, the twin manager **108** includes a device management component **548** for managing devices of a building. The twin manager **108** includes a tenant management component **550** for managing various tenant subscriptions. The twin manager **108** includes an event routing component **552** for routing various events. The twin manager **108** includes an authentication and access component **554** for performing user and/or system authentication and granting the user and/or system access to various spaces, pieces of software, devices, etc. The twin manager **108** includes a commanding component **556** allowing a software application and/or user to send commands to physical devices. The twin manager **108** includes an entitlement component **558** that analyzes the entitlements of a user and/or system and grants the user and/or system abilities based on the entitlements. The twin manager **108** includes a telemetry component **560** that can receive telemetry data from physical systems and/or devices and ingest the telemetry data into the graph **529**. Furthermore, the twin manager **108** includes an integrations component **562** allowing the twin manager **108** to integrate with other applications.

(100) The twin manager **108** includes a gateway **506** and a twin connector **508**. The gateway **506** can be configured to integrate with other systems and the twin connector **508** can be configured to allow the gateway **506** to integrate with the twin manager **108**. The gateway **506** and/or the twin connector **508** can receive an entitlement request **502** and/or an inference request **504**. The entitlement request **502** can be a request received from a system and/or a user requesting that an AI agent action be taken by the AI agent **570**. The entitlement request **502** can be checked against entitlements for the system and/or user to verify that the action requested by the system and/or user is allowed for the user and/or system. The inference request **504** can be a request that the AI agent **570** generates an inference, e.g., a projection of information, a prediction of a future data measurement, an extrapolated data value, etc.

(101) The cloud platform **106** is shown to receive a manual entitlement request **586**. The request **586** can be received from a system, application, and/or user device (e.g., from the applications **110**, the building subsystems **122**, and/or the user device **176**). The manual entitlement request **586** may be a request for the AI agent **570** to perform an action, e.g., an action that the requesting system and/or user has an entitlement for. The cloud platform **106** can receive the manual entitlement request **586** and check the manual entitlement request **586** against an entitlement database **584** storing a set of entitlements to verify that the requesting system has access to the user and/or system. The cloud platform **106**, responsive to the manual entitlement request **586** being approved, can create a job for the AI agent **570** to perform. The created job can be added to a job request topic **580** of a set of topics **578**.

(102) The job request topic **580** can be fed to AI agents **570**. For example, the topics **580** can be fanned out to various AI agents **570** based on the AI agent that each of the topics **580** pertains to (e.g., based on an identifier that identifies an agent and is included in each job of the topic **580**). The AI agents **570** include a service client **572**, a connector **574**, and a model **576**. The model **576** can be loaded into the AI agent **570** from a set of AI models stored in the AI model storage **568**. The AI model storage **568** can store models for making energy load predictions for a building, weather forecasting models for predicting a weather forecast, action/decision models to take certain actions responsive to certain conditions being met, an occupancy model for predicting occupancy of a space and/or a building, etc. The models of the AI model storage **568** can be neural networks (e.g., convolutional neural networks, recurrent neural networks, deep learning networks, etc.), decision trees, support vector machines, and/or any other type of artificial intelligence, machine learning, and/or deep learning category. In some embodiments, the models are rule based triggers and actions that include various parameters for setting a condition and defining an action.

(103) The AI agent **570** can include triggers **595** and actions **597**. The triggers **595** can be conditional rules that, when met, cause one or more of the actions **597**. The triggers **595** can be executed based on information stored in the graph **529** and/or data received from the building subsystems **122**. The actions

597 can be executed to determine commands, actions, and/or outputs. The output of the actions **597** can be stored in the graph **529** and/or communicated to the building subsystems **122**.

(104) The AI agent **570** can include a service client **572** that causes an instance of an AI agent to run. The instance can be hosted by the artificial intelligence service client **588**. The client **588** can cause a client instance **592** to run and communicate with the AI agent **570** via a gateway **590**. The client instance **592** can include a service application **594** that interfaces with a core algorithm **598** via a functional interface **596**. The core algorithm **598** can run the model **576**, e.g., train the model **576** and/or use the model **576** to make inferences and/or predictions.

(105) In some embodiments, the core algorithm **598** can be configured to perform learning based on the graph **529**. In some embodiments, the core algorithm **598** can read and/or analyze the nodes and relationships of the graph **529** to make decisions. In some embodiments, the core algorithm **598** can be configured to use telemetry data (e.g., the timeseries data **564**) from the graph **529** to make inferences on and/or perform model learning. In some embodiments, the result of the inferences can be the timeseries **566**. In some embodiments, the timeseries **564** is an input into the model **576** that predicts the timeseries **566**.

(106) In some embodiments, the core algorithm **598** can generate the timeseries **566** as an inference for a data point, e.g., a prediction of values for the data point at future times. The timeseries **564** may be actual data for the data point. In this regard, the core algorithm **598** can learn and train by comparing the inferred data values against the true data values. In this regard, the model **576** can be trained by the core algorithm **598** to improve the inferences made by the model **576**.

(107) Referring now to FIG. **6**, a process **600** for executing an artificial intelligence agent to infer and/or predict information is shown, according to an exemplary embodiment. The process **600** can be performed by the system **500** and/or components of the system **500**. The process **600** can be performed by the building data platform **100**. Furthermore, the process **600** can be performed by any computing device described herein.

(108) In step **602**, the twin manager **108** receives information from a physical device and stores the information, or a link to the information, in the graph **529**. For example, the telemetry component **560** can receive telemetry data from physical devices, e.g., the building subsystems **122**. The telemetry can be measured data values, a log of historical equipment commands, etc. The telemetry component **560** can store the received information in the graph **529** by relating a node storing the information to a node representing the physical device. For example, the telemetry component **560** can store timeseries data as the timeseries **566** along by identifying that the physical device is a VAV represented by the VAV node **516** and that an edge **537** relates the VAV node **516** to the timeseries node **566**.

(109) In step **604**, the twin manager **108** and/or the cloud platform **106** receives an indication to execute an artificial intelligence agent of an entity represented in the graph **529**, the AI agent being associated with a model. In some embodiments, the indication is created by a user and provided via the user device **176**. In some embodiments, the indication is created by an application, e.g., one of the applications **110**. In some embodiments, the indication is a triggering event that triggers the agent and is received from the building subsystems **122** and/or another agent (e.g., an output of one agent fed into another agent).

(110) In some embodiments, the AI agent is an agent for a specific entity represented in the graph **529**. For example, the agent could be a VAV maintenance agent configured to identify whether a VAV (e.g., a VAV represented by the nodes **512**, **520**, and/or **516**) should have maintenance performed at a specific time. Another agent could be a floor occupant prediction agent that is configured to predict the occupancy of a particular floor of a building, e.g., the floor represented by the floor node **522**.

(111) Responsive to receiving the indication, in step **606**, the AI agent **570** causes a client instance **592** to run the model **576** based on the information received in step **602**. In some embodiments, the information received in step **602** is provided directly to the AI agent **570**. In some embodiments, the information is read from the graph **529** by the AI agent **570**.

(112) In step **608**, the AI agent **570** stores the inferred and/or predicted information in the graph **529** (or stores the inferred and/or predicted information in a separate data structure with a link to the graph **529**). In some embodiments, the AI agent **570** identifies that the node that represents the physical entity that the AI agent **570** inferred and/or predicted information for, e.g., the VAV represented by the VAV **516**. The AI agent **570** can identify that the timeseries node **566** stores the inferred and/or predicted information by identifying the edge **537** between the VAV node **516** and the timeseries node **566**.

(113) In step **610**, the AI agent **570** can retrieve the inferred or predicted information from the graph **529** responsive to receiving an indication to execute the model of the AI agent **570** of the inferred or predicted information, e.g., similar to the step **604**. In step **612**, the AI agent **570** can execute one or more actions based on the inferred and/or predicted information of the step **610** based the inferred and/or predicted information retrieved from the graph **529**. In some embodiments, the AI agent **570** executes the model **576** based on the inferred and/or predicted information.

(114) In step **614**, the AI agent **570** can train the model **576** based on the inferred or predicted information read from the graph **529** and received actual values for the inferred or predicted information. In some embodiments, the AI agent **570** can train and update parameters of the model **576**. For example, the timeseries **564** may represent actual values for a data point of the VAV represented by the VAV node **516**. The timeseries **566** can be the inferred and/or predicted information. The AI agent **570** can compare the timeseries **564** and the timeseries **566** to determine an error in the inferences and/or predictions of the model **576**. The error can be used by the model **576** to update and train the model **576**.

(115) Referring now to FIG. **7**, a process **700** of an agent executing a trigger rule and an action rule is shown, according to an exemplary embodiment. The process **700** can be performed by the system **500** and/or components of the system **500**. In some embodiments, the building data platform **100** can perform the process **700**. Furthermore, the process **700** can be performed by any computing device described herein.

(116) In step **702**, the building data platform can store an agent **570** in a data structure. The agent **570** can include a trigger rule indicating a condition for executing an action rule and an action rule indicating an action to be performed responsive to the condition being met. In some embodiments, the model **576** includes, or can be replaced with, the trigger rule and the action rule. The trigger rule and the action rule can be logical statements and/or conditions that include parameter values and/or create an output action.

(117) In step **704**, the agent **570** can receive information from at least one of a physical device and/or from the graph **529**. The information can be generated by a physical device, e.g., the building subsystems **122**. The building data platform **100** can, in some embodiments, receive the information from the physical device, ingest the information into the graph **529**, and the agent **570** can read the information from the graph **529**. In some embodiments, the agent **570** can check the information of the graph **529** against a trigger rule at a set period.

(118) In step **706**, the agent **570** determines whether the information received in the step **704** causes the condition to be met. The agent **570** can apply the information to the trigger rule to determine whether the trigger rule is triggered, i.e., the condition of the trigger rule being met.

(119) In step **708**, the agent **570** can perform the action responsive to the condition being met by the information determined in step **706**. The action may cause a physical device to be operated or information be sent to another agent including another trigger rule and another action rule. In some embodiments, the action can be performed by executing the action rule of the agent **570**. The action rule can perform an action based on one or more parameter value of the action rule. In some embodiments, the action output of the action rule can be sent directly to the physical device, e.g., the building subsystems **122**. In some embodiments, the action output can be stored into the graph **529**. Another operating component of the building data platform **100**, e.g., the command processor **136**, can read the action from the graph **529** can communicate a corresponding command to the building subsystems **122**.

(120) Referring now to FIG. **8**, a system **800** where a clean air optimization (CAO) AI service **804** and an energy prediction model (EPM) AI service **806** operate to make inferences and predictions based on data of a knowledge graph **802** is shown, according to an exemplary embodiment. The knowledge graph **802** includes various nodes and edges. The nodes may be the nodes **808-842**. The edges may be the edges **844-878**.

(121) The nodes may represent various entities of a building and/or buildings. The entities may be a campus, a building, a floor, a space, a zone, a piece of equipment, a person, a control point, a data measurement point, a sensor, an actuator, telemetry data, a piece of timeseries data, etc. The edges **844-878** can interrelate the nodes **808-842** to represent the relationships between the various entities of the building. The edges **844-878** can be semantic language based edges **844-878**. The edges can include words and/or phrases that represent the relationship. The words and/or phrases can include at least one predicate, in some cases.

(122) The knowledge graph **802** includes a building node **836** representing a building. The building can

include floors, represented by the building node **836** being related to a floor **1** node **840** via an “hasPart” edge **870** and the building node **836** being related to a floor **3** node **838** via a “feeds” edge **868**. The building includes an energy prediction, e.g., a value or a timeseries of values indicating energy usage of the building. This can be represented by the building node **836** being related to an energy prediction node **830** via the edge **872**.

(123) The floor **1** includes zones, indicated by the floor **1** node **840** being related to a zone**1-1** node **808** via a “hasPart” edge **844** and the floor **1** node **840** being related to a zone**1-2** node **842** via a “hasPart” edge **845**. Furthermore, the floor **1** can be fed by a particular AHU, AHU**1**. This is indicated by an AHU**1** node **822** being related to the floor **1** node **840** via the “feeds” edge **860**. The zone**1-2** can include a VAV**1-2** that feeds air to it. This can be indicated by the VAV**1-2** node **810** being related to the zone**1-2** node **842** by the “feeds” node **862**. The AHU**1** can feed the VAV**1-2**, indicated by the AHU**1** node **822** being related to the VAV**1-2** node **810** by the “feeds” edge **858**.

(124) An AHU **3** can feed air to the floor **3**, indicated by the “AHU**3**” node **834** being related to the “floor **3**” node **838** by the “feeds” edge **866**. The “AHU**3**” node **834** is related to a CAO inference **828** via an edge **874** representing that the AHU has a clean air optimization inference that could be determined by the CAO AI service **804**. The knowledge graph **802** includes an AHU**2** node **832**. The AHU**2** **832** is related to a floor **2** node **841** via a “feeds” edge **864** indicating that the AHU**2** feeds air to the floor **2**. The AHU**2** node **832** is related to a CAO inference **826** via the edge **876**. The CAO inference **826** can indicate an inference made by the CAO AI service **804** for the AHU**2**. The AHU**1** **822** is related to a CAO inference **824** via edge **878**. The CAO inference **824** can indicate a clean air optimization inference made by the CAO AI service **804** for the AHU**1**.

(125) The knowledge graph **802** includes a VAV**1-1** node **812** indicating a VAV**1-1**. The VAV**1-1** node **812** is related to the zone**1-1** node **808** via a “feeds” edge **846** indicating that the VAV**1-1** feeds air to the zone**1-1**. The AHU**1** can feed air to the VAV**1-1** indicated by the AHU**1** node **822** being related to the VAV**1-1** node **812** via the “feeds” edge **856**. The VAV**1-1** node **812** includes various points, e.g., a zone temperature point (represented by the VAV**1**.ZN-T node **814**) and a zone relative humidity point (represented by the VAV**1**.ZN-RH node **816**). The VAV**1-1** node **812** is related to the VAV**1**.ZN-T node **814** via the “hasPoint” edge **848**. The VAV**1-1** node **812** is related to the VAV**1**.ZN-RH node **816** via the “hasPoint” edge **854**.

(126) The VAV**1**.ZN-T point includes a timeseries node **818** representing and/or storing a timeseries for the zone temperature, indicated by the timeseries node **818** being related to the VAV**1**.ZN-T node **814** via the “hasTimeseriesID” node **850**. The VAV**1**.ZN-RH point includes a timeseries node **820** representing and/or storing a timeseries for the zone humidity, indicated by the timeseries node **820** being related to the VAV**1**.ZN-RH node **816** via the “hasTimeseriesID” node **852**. In some embodiments, the timeseries node **818** and the timeseries node **820** are identifiers of a particular timeseries stored in a separate timeseries database, the identifier uniquely identifying the location of the timeseries so a system can consult the knowledge graph **802** and use the identifiers to retrieve the timeseries data from the separate timeseries database.

(127) The system **800** includes a CAO AI service **804**. The CAO AI service **804** can be configured to identify timeseries that it needs to execute on. For example, if the CAO AI service **804** is executing for the AHU**1**, the CAO AI service **804** could identify timeseries data linked to the AHU**1**. The CAO AI service **804** can generate CAO inferences, e.g., can infer ideal settings for clean air. The ideal settings could be an AHU supply air temperature setpoint, an AHU minimum ventilation rate, etc. The ideal settings can be ingested into the knowledge graph **802**, e.g., as the CAO inferences **824-830**.

(128) In some embodiments, the CAO AI service **804** (or an agent for the CAO AI service **804**) operates on behalf of a particular entity, e.g., the AHU **1**. The CAO AI service **804** can generate inferences with data of the AHU **1**, e.g., by identifying timeseries data of the AHU **1** by identifying timeseries nodes of the knowledge graph **802** via an edge. The inferences can be ingested into the knowledge graph **802** by generating a new node and/or adding an edge between the new node and the node of the AHU **1**, AHU **1** node **822**. Similarly, the inferences can be added to the knowledge graph **802** by updating an existing node related to the AHU **1** node **822** via an existing edge. In some embodiments, the inferences of the CAO AI service **804** can generate a recommendation, e.g., a control setting for improving or optimizing air quality, which can be reviewed and approved by a user via the user device **176**.

(129) The EPM AI service **806** can generate energy predictions for various buildings, spaces, or devices of

a building, e.g., entitles of the knowledge graph **802**. For example, the EPM AI service **806** could predict a future energy consumption level of the building **836**, e.g., a future energy demand. The energy prediction can be a node **830** related to the building node **836** via the edge **872**. In some embodiments, the EPM AI service **806** can generate the energy prediction node **830** responsive to generating the future building energy consumption and cause the node **830** to include a value for the future building energy consumption. The node **830** can be added to the graph **802** and the edge **872** can be generated by the EPM AI service **806** and added between the building node **836** and the energy prediction **830**.

(130) In some embodiments, the energy prediction node **830** already exists within the knowledge graph **802**. In this example, the EPM AI service **806** can identify the energy prediction node **830** by identifying an edge **872** between the building **836** and the energy prediction node **830**. The EPM AI service **806** can then ingest the energy prediction into the node **830**.

(131) Referring now to FIG. **9**, a system **900** including a CAO AI agent **924** and an EPM AI agent **926** operating against the knowledge graph **802** to generate inferences and predictions is shown, according to an exemplary embodiment. The system **900** can be implemented on one or more processing circuits, e.g., as instructions stored on one or more memory devices and executed on one or more processors. The memory devices and processors may be the same as or similar to the memory devices and processors described with reference to FIG. **1**.

(132) The CAO AI agent **924** can operate on behalf of the CAO AI service **804**. Similarly, the EPM AI agent **926** can operate on behalf of the EPM AI service **806**. Furthermore a service bus **914** can interface with the agent **924** and/or the agent **926**. A user can interface with the agents **924-926** via the user device **176**. The user can provide an entitlement request, e.g., a request that the user is entitled to make and can be verified by the AI agent manager **916**. The AI agent manager **916** can send an AI job request based on a schedule to the service bus **914** based on the entitlement request. The service bus **914** can communicate the AI job request to the appropriate agent and/or communicate results for the AI job back to the user device **176**.

(133) In some embodiments, the CAO AI agent **924** can provide a request for generating an inference to the CAO AI service **804**. The request can include data read from the knowledge graph **802**, in some embodiments.

(134) The CAO AI agent **924** includes a client **902**, a schema translator **904**, and a CAO client **906**. The client **902** can be configured to interface with the knowledge graph **802**, e.g., read data out of the knowledge graph **802**. The client **902** can further ingest inferences back into the knowledge graph **802**. For example, the client **902** could identify timeseries nodes related to the AHU1 node **2622**, e.g., timeseries nodes related to the AHU1 node **2622** via one or more edges. The client **902** can then ingest the inference made by the CAO AI agent **924** into the knowledge graph **802**.

(135) The client **902** can provide data it reads from the knowledge graph **802** to a schema translator **904** that may translate the data into a specific format in a specific schema that is appropriate for consumption by the CAO client **906** and/or the CAO AI service **804**. The CAO client **906** can run one or more algorithms, software components, machine learning models, etc. to generate the inference and provide the inference to the client **902**. In some embodiments, the client **906** can interface with the EPM AI service **806** and provide the translated data to the EPM AI service **806** for generating an inference. The inference can be returned by the EPM AI service **806** to the CAO client **906**.

(136) The EPM AI agent **926** can operate in a similar manner to the CAO AI agent **924**, in some embodiments. The client **908** can retrieve data from the knowledge graph **802** and provide the data to the schema translator **910**. The schema translator **910** can translate the data into a readable format by the CAO AI service **804** and can provide the data to the EPM client **912**. The EPM client **912** can provide the data along with a prediction request to the CAO AI service **804**. The CAO AI service **804** can generate the prediction and provide the prediction to the EPM client **912**. The EPM client **912** can provide the prediction to the client **908** and the client **908** can ingest the prediction into the knowledge graph **802**.

(137) In some embodiments, the knowledge graph **802** includes data necessary for the inferences and/or predictions that the agents **924** and **926** generate. For example, the knowledge graph **802** can store information such as the size of a building, the number of floors of the building, the equipment of each floor of the building, the square footage of each floor, square footage of each zone, ceiling heights, etc. The data can be stored as nodes in the knowledge graph **802** representing the physical characteristics of the building. In some embodiments, the EPM AI agent **926** makes the predictions based on the characteristic

data of the building and/or physical areas of the building.

(138) Referring now to FIG. 10, a system **1000** including the edge platform **102** performing event enrichment at the edge platform **102** before the events are communicated to the cloud platform **106** is shown, according to an exemplary embodiment. The system **1000** includes the building subsystems **122**, the edge platform **102**, the cloud platform **106**, the applications **110**, and the twin manager **108**. The edge platform **102** can receive events from the building subsystems **122** and enrich the events before passing the events on to the cloud platform **106**. Because the edge platform **102** is located on-premises, e.g., on the edge, the events can be enriched before being passed on to other cloud systems and/or used in edge based analytics run on the edge platform **102**. In some embodiments, processors, memory devices, and/or networking devices of the edge platform **102** are located on-premises within a building.

(139) The edge platform **102** can receive events from the building subsystems **122**. The events can be data packages describing an event that has occurred with a timestamp of when the event occurred. The events can be raw events that are composed of content that is emitted from a producing system. However, the event may not include any intent or knowledge of the system that consumes it. The event can be of a particular event type. An enrichment manager **1002** of the edge platform **102** can receive the events from the building subsystems **122**. The enrichment manager **1002** can be the same as, or similar to, the enrichment manager **138**.

(140) The enrichment manager **1002** can enrich the events received from the building subsystems **122** based on event context received and/or retrieved from a lite digital twin **1008** of the edge platform **102**. For example, the enrichment manager **1002** can add entity and/or entity relationship information associated with the event to the event to generate the enriched events **1004**. The event enrichment can be the same as or similar to the enrichment described with referenced to FIGS. 1-3 and FIG. 8. The enriched events **1004** can be an event with additional added properties or attributes that provide context regarding the event.

(141) In some embodiments, the enrichment manager **1002** includes multiple event streams. The event streams can be data enrichment processing streams for particular events and/or particular types of events. Each event stream can be linked to a tenant and/or tenant subscription. Each event stream can indicate one or more rules for enriching an event, e.g., an indication of the information to add to the event. In this regard, one event can be applied to multiple event streams and receive different enrichments to generate multiple enriched events. Each enriched event can be provided to a different application or end system.

(142) The edge platform **102** includes edge applications **1010**. The edge applications **1010** can be similar to or the same as the applications **110**. While the applications **110** may be run on a cloud system, the edge applications **1010** can be run locally on the edge platform **102**. The edge applications **1010** can operate based on the enriched events **1004** and may not need to consult a digital twin to acquire context regarding an event since the enriched events **1004** may already include the needed context. In some embodiments, the edge application **1010** performs analytics (e.g., aggregation, data monitoring, etc.), control algorithms, etc. for the building subsystems **122**.

(143) For example the edge applications **1010** can generate control decisions for the building subsystems **122** based on the enriched events **1004**, e.g., temperature setpoints for zones, fan speed settings for fans, duct pressure setpoints, ventilation commands, etc. In some embodiments, the edge applications **1010** include models, e.g., machine learning models for predicting characteristics and/or conditions and/or for operating the building subsystems **122**. In some embodiments, the machine learning is performed at the edge platform **102** which results in higher scores than machine learning performed in the cloud since a greater amount of data can be collected faster and used for training at the edge.

(144) In some embodiments, the enrichment manager **1002** only operates when the twin manager **108** is not operating and enriching events. For example, the edge platform **102** can receive an indication that there is an error with cloud systems, e.g., network issues, computing issues, etc. In this regard, the edge platform **102** can take over enriching the events with the enrichment manager **1002** and operating on the events with the edge applications **1010**. In this regard, the enrichment and application operation can dynamically move between the edge platform **102** and the cloud platform **106**. Furthermore, load balancing can be implemented so that some events are enriched and operated on by edge applications **1010** while other events are passed to the cloud platform **106** and/or the twin manager **108** for enrichment and provided to the applications **110** for operation.

(145) In some embodiments, by performing enrichment at the edge platform **102**, analytics can be

performed at the edge platform **102** based on the enriched events. In this regard, lower latencies can be realized since analytics and/or control algorithms can be performed quickly at the edge platform **102** and data does not need to be communicated to the cloud. In some embodiments, the edge applications **1010** and/or machine learning models of the edge applications **1010** can be built in the cloud and communicated to the edge platform **102** and additional learning can be performed at the edge platform **102**.

(146) The edge platform **102** includes the lite digital twin **1008**. The lite digital twin **1008** can be a version of a digital twin **1011** of the twin manager **108**. The digital twins **1011** and/or **1008** can be virtual representations of a building and/or the building subsystems **122** of the building. The digital twin **1011** and/or the digital twin **1008** can be or can include the graph projection database **162**, e.g., one or more graph data structures. The digital twin **1011** and/or the lite digital twin **1008** can be the graphs shown in FIGS. **11-13**. In some embodiments, the lite digital twin **1008** is a projection that does not include all nodes and edges of a full projection graph. The lite digital twin **1008** may only include the nodes or edges necessary for enriching the events and can be built on projection rules that define the information needed that will be used to enrich the events.

(147) In some embodiments, the lite digital twin **1008** can be synchronized, in whole or in part, with the digital twin **1011**. The lite digital twin **1008** can include less information than the digital twin **1011**, e.g., less nodes or edges. The lite digital twin **1008** may only include the nodes and/or edges necessary for enriching events of the building subsystems **122**. In some embodiments, changes or updates to the digital twin **1011** can be synchronized to the lite digital twin **1008** through a change feed of change feed events. The change feed can indicate additions, removals, and/or reconfigurations of nodes or edges to the graph projection database **162**. Each change feed event can indicate one update to the digital twin **1011**.

(148) A digital twin updater **1006** can receive the events of the change feed from the change feed generator **152** and update the lite digital twin **1008** based on each change feed event. The updates made to the lite digital twin **1008** can be the same updates as indicated by the events of the change feed. In some embodiments, the digital twin updater **1006** can update the lite digital twin **1008** to only include the nodes and edges necessary for enrichment of the events, and thus include less nodes and edges than the digital twin **1011**.

(149) In some embodiments, the digital twin updater **1006** filters out change feed events if the change feed events do not pertain to information needed to enrich the events. In this regard, the digital twin updater **1006** can store a list of information needed for enrichment, e.g., the digital twin updater **1006** can include all event subscriptions or enrichment rules. The digital twin updater **1006** can determine whether a change feed event updates information pertaining to event enrichment and only update the lite digital twin **1008** responsive to determining that the change feed event updates information needed for enrichment. In some embodiments, when a new event subscription and/or new enrichment rule is created, the digital twin updater **1006** can communicate with the digital twin **1011** to retrieve nodes or edges needed for the new event subscription and/or enrichment rules.

(150) Referring now to FIG. **11**, a building graph **1100** divided and distributed across a cloud system and an edge system is shown, according to an exemplary embodiment. The building graph **1100** can form a hybrid knowledge graph (e.g., a graph data structure) which may be part of, or form, a digital twin. The hybridization of the building graph **1100** can allow for various applications to run against the building graph **1100** on edge systems, in cloud systems, or in a combination of edge systems and cloud systems. In some embodiments, the building graph **1100** can be divided up based on space or building. For example, edge systems of one space might have a building graph for the one space while edge systems for another space might have a building graph for the other space. In some embodiments, the building graph **1100** is divided up based on the operational requirements of various edge systems and cloud systems, e.g., a portion of the building graph **1100** may be stored and operated on each edge system or cloud system that allows each edge system or cloud system to perform its operations.

(151) The building graph **1100** includes a cloud building graph **1102** and an edge building graph **1104**. The cloud building graph **1102** can be stored by a cloud system, in some embodiments. The edge building graph **1104** can be stored by an edge system, in some embodiments, examples of the cloud system and the edge system are provided in FIG. **12**. The building graphs **1102** and **1104** can be the same as, or similar to, the building graphs described in FIGS. **1-5** and **8**.

(152) By dividing the building graph **1100** into parts (e.g., the graph **1102** and the graph **1104**) and distributing the parts across various locations provides privacy and security advantages. Because the

building graph **1100** is hacked and distributed, for a hacker or other entity to gain access to the entire building graph **1100**, multiple separate devices (and the security systems of each device) may need to first be compromised. Requiring multiple different devices to be compromised before the entire building graph **1100** is compromised can increase the difficulty of compromising the building graph **1100**. Furthermore, by distributing parts of the building graph **1100**, privacy can be established. In some embodiments, the information of each graph being stored separately can allow for sensitive information to reside in one device with less sensitive information residing in a different device. This allows for the location of information to be controlled to make sure that private information is only stored in locations that a user or entity desires (e.g., locations with high security).

(153) The cloud building graph **1102** and the edge building **1104** include various nodes, i.e., nodes **1106-1134**. The cloud building graphs **1102** and **1104** further includes edges connecting the various nodes **1106-1134**, e.g., edges **1136-1174**. The nodes **1106-1134** can represent buildings, spaces, systems, devices, people, state data, timeseries data, sensor measurements, event data, inferred information, event streams, and/or any other type of data. The edges **1136-1174** can represent directional or bidirectional relationships between the nodes **1106-1134**. In some embodiments, the edges **1136-1174** include semantic language based relationships that include one or more predicates, nouns, and/or prepositions that describe a relationship type. In some embodiments, the cloud building graph **1102** and/or the edge building graph **1104** are formatted according to a specific schema, e.g., the BRICK schema. In some embodiments, the nodes and edges stored by the cloud building graph **1102** and the edge building graph **1104** are different. In some embodiments, there is overlap between the nodes and edges stored by the cloud building graph **1102** and the edge building graph **1104**.

(154) In some embodiments, the cloud building graph **1102** includes a building node **1122**. The building node **1122** can represent a particular building, e.g., a school, a hospital, a commercial business office, a high rise, an apartment, etc. The cloud building graph **1102** includes a floor 2 node **1124** related to the building node **1122** by a “hasPart” node **1158**. This indicates that the building has a second floor, floor 2. The building node **1122** is related to a floor 3 node **1120** by a “hasPart” edge **1156**. This indicates that the building includes a third floor, floor 3. Furthermore, the building node **1122** is related to a floor 1 node **1118** by a “hasPart” edge **1152**. This indicates that the building has a first floor, floor 1.

(155) The floor 1 node **1118** is related to a zone1-1 node **1116** by a “hasPart” edge **1150**. This indicates that the first floor, floor 1, has a first zone, zone1-1. The zone1-1 node **1116** is related to a VAV1-2 node **1114** by a “feeds” edge **1148**. This indicates that a particular VAV, VAV1-2, operates to feed air to the zone1-1. The VAV1-2 is associated with two different points and timeseries. The VAV1-2 node **1114** is related by a “hasPoint” edge **1140** to a zone temperature point, VAV1.ZN.T node **1110**. The VAV1.ZN.T node **1110** is related by a “hasTimeseries” edge **1136** to a timeseries node **1106**. This indicates that the zone temperature point includes an associated timeseries of zone temperature measurements made by the VAV1-2 or another system associated with the zone1-1 (e.g., a sensor or thermostat). The VAV1-2 node **1114** is related by a “hasPoint” edge **1142** to a zone humidity point, VAV1.ZN.RH node **1112**. The VAV1.ZN.RH node **1112** is related by a “hasTimeseries” edge **1138** to a timeseries node **1108**. This indicates that the zone humidity point includes an associated timeseries of zone humidity measurements made by the VAV1-2 or another system associated with the zone1-1 (e.g., a sensor or thermostat). In some embodiments, the timeseries node **1106** and/or the timeseries node **1108** (or the timeseries data associated with the nodes) are stored in a separate timeseries database and related to the graph **1102** by an identifier.

(156) The edge building graph **1104** includes an AHU1 node **1128** representing an AHU of the building. The AHU1 node **1128** is related to a VAV1-2 node **1130** by a “feeds” edge **1170**. This indicates that the AHU1 feeds air to the VAV1-2. The VAV1-2 node **1130** is related to a zone1-2 node **1132** by a “feeds” edge **1172**. This indicates that the VAV1-2 node **1130** feeds air to the zone1-2. The AHU1 node **1128** is related by a “hasA” edge **1174** to an artificial intelligence (AI) inference event stream node **1134**. The AI inference event stream node **1134** can store (or include an identifier linking to a separate data storage system) inferred events generated by an AI service. The event stream could be inferred weather conditions, inferred occupancy, inferred comfort settings for the AHU, etc. A system that runs the AHU1 can use the event stream to make operational decisions (e.g., determine heating levels, cooling levels, air pressure settings, fan levels, air mixture levels, etc.).

(157) The cloud building graph **1102** and the edge building graph **1104** may both describe elements of the same building, in some embodiments. Therefore, the cloud building graph **1102** and the edge building

graph **1104** can be related together. The interrelationships can enable the cloud building graph **1100** to be split apart but for queries or other actions to still consider the entire graph **1100**, in some embodiments. (158) In some embodiments, an edge **1** node **1126** can be stored by the cloud building graph **1102** and/or the edge building graph **1104**. The edge **1** node **1126** may represent the edge building graph **1104**. Similarly, another node might describe the cloud building graph **1102** (e.g., as shown in FIG. 12). The graphs **1102** and/or **1104** can include edges that relate nodes of each graph **1102** and/or **1104** to the node **1126**. Relating a node to the node **1126** may indicate that an object of the relationship exists in another graph, e.g., the graph indicated by the edge **1** node **1126**. For example, if the edges are in the form of triples with subjects, predicates, and objects, relating a node (subject) via an edge (predicate) to the edge **1** node **1126** may indicate that the object for the relationship exists in another graph (e.g., a graph indicated by the node **1126**). Similarly, if the edge **1** node **1126** is the subject of a triple, this may indicate that the subject of the triple is stored in another graph, e.g., the graph indicated by the edge **1** node **1126**.

(159) The VAV1-2 node **1114** is related to the AHU1 node **1128** via an edge “Edge1[AHU1].fwdarw.feeds::Cloud1[VAV1-2]” **1144** between the edge **1** node **1126** and the VAV1-2 node **1114** and the edge “Edge1[AHU1].fwdarw.feeds::Cloud1[VAV1-2]” edge **1164**. The edges **1144** and **1164** can indicate that the AHU1 node **1128** (the subject) feeds (the predicate) the VAV1-2 node **1114** (object). The edges **1144** and **1164** can each identify the AHU1 node by name, provide an indication of which graph the AHU1 node is stored, e.g., “Edge1[AHU1]” provide an indication of the direction of the relationship “.fwdarw.” (from the AHU1 to the VAV1-2) and indicate the location of the VAV1-2 node **1114** “Cloud1[VAV1-2].” The “::” symbol can indicate a separation between graphs, e.g., a separation between the Edge1 graph (the graph **1104**) and the Cloud1 graph (the graph **1102**). The location of the “AHU1” on the left side of the arrow.fwdarw.can indicate that the AHU1 is the subject of the triple while the location of the VAV1-2 on the right side of the arrow “.fwdarw.” can indicate that the VAV1-2 is the object of the triple. Reversing the arrow “.fwdarw.” to “<-” would indicate a reverse of the subject and object, in some embodiments. Furthermore, an arrow such as “<.fwdarw.” would indicate a bidirectional relationship where each node is both a subject and object, in some embodiments. The “feeds” portion of the edges **1144** and **1164** can indicate the semantic type (e.g., “isPartOf,” “includes,” etc.) of the edges **1144** and **1164**.

(160) While a particular syntax is provided in FIG. 11 for the edges relating nodes of different graphs, various other syntaxes could be used to relate the information. Furthermore, the relation of information could, in some embodiments, be stored within the nodes or edges of the graph **1100** itself and/or within a separate lookup table. In some embodiments, the lookup table could relate the graph names “Cloud1” and/or “Edge1” to physical devices that store the respective graphs, e.g., to system identifiers, to network addresses, to cryptographic communication codes, etc. In some embodiments, the lookup table can use obfuscation to hide identifying information of the devices storing the graphs **1102** and **1104**.

(161) The “Edge1[AHU1].fwdarw.feeds Cloud[Floor1]” edge **1146** between the edge **1** node **1126** and the floor **1** node **118** and the “Edge1[AHU1].fwdarw.feeds Cloud1[Floor1]” edge **1166** between the AHU1 node **1128** and the edge1 node **1126** can indicate that the AHU 1 feeds air to the floor 1. The “Cloud1[Floor1].fwdarw.hasPart Edge1[Zone1-2]” edge **1154** between the floor **1** node **1118** and the edge **1** node **1126** and the “Cloud1[Floor1].fwdarw.hasPart Edge1[Zone1-2]” edge **1168** between the edge **1** node **1126** and the zone1-2 node **1132** can indicate that the floor 1 has the zone1-2. The “Edge1[AHU1].fwdarw.feeds Cloud1[Floor2]” edge **1160** between the edge **1** node **1126** and the floor 2 node **1124** and the “Edge1[AHU1].fwdarw.feeds Cloud1[Floor2]” edge **1162** between the edge **1** node **1126** and the AHU 1 node **1128** can indicate that the AHU1 feeds the floor 2.

(162) By dividing the building graph **1100** into the graphs **1102** and **1104** but including the edges **1144-1146, 1154, 1160, 1162, 1164, 1166, and 1168** which link the graphs **1102** and **1104** together, queries can be executed which traverse the entire building graph **1100** without requiring the queries to specify which graph is being queried. For example, if a system provides a query for all parts of the floor 1, the query response can include the zone1-1 and the zone1-2, even though the zone 1-2 is stored in a graph separate from the floor 1. The query responding system can identify and retrieve an indication of the zone 1-2 based on traversing edges across the graphs **1102** and **1104**, e.g., the edges **1154** and **1168**. In some embodiments, a first system that is associated with the cloud building graph **1102** (e.g., stores the graph **1102**) can complete a first portion of the query (e.g., identify the zone1-1) use the edge **1154** to identify that a second result for the query is stored in the graph **1104**, and second the query to a second system

associated with the edge building graph **1104**. The second system can complete the query (e.g., identify the zone **1-2** node **1132** based on the edge **1168** or any other pieces of information in the graph **1104**), and send a response to the first system.

(163) The first system can then combine the results together and generate a query response based on the combined results. One or more application programming interfaces (APIs) can provide an interface for querying the graphs **1102-1104**. Because the graphs **1102** and **1104** are linked together via edges, even though they may be stored in separate data storage devices, on different devices connected through one or more networks, etc., the queries provided to the API can be agnostic to (e.g., not include) information identifying where (or in what graph) the queried information resides, how to break down the queries into sub-parts, and/or how to route the queries. The API can break a query into various parts for each of the graphs **1102-1104** and route the queries to various devices to query the graphs stored by each device. One example of breaking a query down into sub-parts to produce a query result is found in U.S. patent application Ser. No. 15/409,489 filed Jan. 18, 2017 (now U.S. Pat. No. 10,480,804), the entirety of which is incorporated by reference herein.

(164) In some embodiments, instead of including timeseries nodes, e.g., the node **1110** and/or the node **1112**, the graph **1100** can store events, timeseries, and/or other data samples linked directly with an entity node (e.g., equipment, space, person, etc.). For example, a single edge may connect the data sample node directly to the entity node. In some embodiments, each entity, whether equipment, space, or a point, can include an edge to an event metadata node. The event metadata node can describe the type of data samples that can be associated with that entity. The data samples can be formatted according to a format defined by the metadata node but linked directly with the entity generating them, or in the case of inferred events, applied to them.

(165) For example, if there is a meeting that gets booked in a conference room, the event describing that meeting room booking would be associated directly with the conference room entity itself. Furthermore, an event metadata entity can also be linked to a node representing the meeting room which describes a “MeetingRequest” event and its format. This metadata entity may describe the types of events that can be associated with that conference room entity.

(166) Referring now to FIG. **12**, a system **1200** is shown including the twin manager **108** and an edge device **1204** storing the building graph **1100** divided and distributed across the twin manager **108** and the edge device **1204**, according to an exemplary embodiment. The twin manager **108** can be a cloud based device that includes the memory device(s) **150** and the processor(s) **148**. The one or more memory device(s) **150** can store the cloud building graph **1102**, the cloud agents **1206**, the device-graph lookup table **1220**, and/or any instructions that can be executed by the processor(s) **140** to store, manage, execute against, etc. the cloud building graph **1102** and/or the edge building graph **1104**.

(167) The edge device **1204** includes processor(s) **1210** and the memory device(s) **1212**. The processor(s) **1210** and the memory device(s) **1212** can be similar to, or be the same type as, the processor(s) **148** and the memory device(s) **150**. The edge device **1204** can be a device located on the edge, e.g., on-premises within a building, at a particular environment, in a vehicle, etc. The edge device **1204** could be a building controller, a vehicle computing system, a personal computer, an equipment controller, a METASYS ADX server, a building server, a fire response system, an access control system, a video surveillance system, etc. In some embodiments, the edge device **1204** can be a server of a cellular tower. For example, the edge device **1204** could be a 5G Mobile Edge Computing (MEC) system. 5G MEC systems can allow for communication with edge devices and processing for edge devices while only requiring communication via 5G, i.e., without requiring any other network hops (e.g., network hops across a core network, the Internet, etc.). 5G MEC is described in U.S. patent Ser. No. 17/221,064 filed Apr. 2, 2021, the entirety of which is incorporated by reference herein.

(168) As an example, in some embodiments, the edge device **1204** is a field controller or an equipment controller. The controller could collect and store historical equipment data which can be stored in, or linked to, a building graph stored by the controller. The controller can further run various agents to perform fault detection and diagnosis for the equipment. The agents can run against the historical data, the graph stored by the controller, or another graph stored in a cloud or other edge system that includes information relating to making a fault detection for the equipment. For example, the other graphs could store information about the spaces, people, or equipment that might provide information that the controller could use in making a fault detection and/or diagnosis. For example, the information might indicate how

other similar pieces of equipment are operating, the environmental data of various spaces controlled by the equipment, etc.

(169) In some embodiments, the distributed digital twin can store both the portion of the building graph and the historical data necessary for the agents of the digital twin to execute. For example, the edge device **1204** could store a historical database, e.g., a timeseries database. In some embodiments, the timeseries data can be stored in the timeseries database with a link to a node of the edge building graph **1104**. In this regard, the edge device **1204** could receive sensor measurements from the sensor devices **1214** and/or actuator feedback from the actuator devices **1216** and store the data, via a device manager **1218** of the edge device **1204**, in the graph **1104** and/or in a historical database of the edge device **1204**. In some embodiments, the device manager **1218** can store a link in a node of the edge building graph **1104** that links to a location in a separate database that aggregates the received sensor and/or actuator data.

(170) In some embodiments, for a campus of multiple buildings or an entity with multiple building sites, one server could reside at each building. Each of the servers could store a building graph for the specific building that the servers are located in. The cloud server could store a backup graph of all of the graphs stored at each of the individual servers. In some embodiments, the cloud server could include a graph storing other high level information that is related to the information of one or more of the graphs of the servers. The servers could, in some embodiments, push gathered building data up to the cloud for the cloud to run algorithms against. In some embodiments, the servers could retain the gathered building data and run algorithms for the respective buildings locally.

(171) The twin manager **108** includes a device-graph lookup table while the edge device **1204** includes a device-graph lookup table **1222**. The tables **1220** and/or **1222** can be tables that link a graph name to a specific device. The tables **1222** and/or **1220** can allow the interface **1224** of the twin manager **108** and/or the interface **1226** of the edge device **1204** to communicate with each other via the network **104**. For example, the communication may be querying the graphs **1102** and/or **1104**. The interfaces **1224** and/or **1226** can be circuits and/or pieces of software (e.g., API code) that handles communication, queries, commands, questions, etc. between the twin manager **108** and the edge device **1204**.

(172) In some cases, the interfaces **1224** and/or **1226** are shepherding agents that coordinate the local storage of a graph and/or agent between multiple devices. In some embodiments, the shepherding agents can coordinate communication between edge and cloud systems such that a gateway layer (e.g., the edge platform **102**) is not needed to facilitate communication between the cloud (e.g., the cloud platform **106** and/or the twin manager **108**) and edge devices (e.g., the building subsystems **122**). The messaging and/or routing mechanisms provided by the interfaces **1224** and/or **1226** can allow for business logic, control algorithms, artificial intelligence (AI) models, machine learning (ML) models, etc. that might normally run in the cloud to be run in various edge or cloud systems.

(173) The device-graph lookup table **1220** and/or the device-graph lookup table **1222** can link graph names of the graphs **1102** and/or **1104** to device identifiers and/or communication addresses of the twin manager **108** and/or edge device **1204**. An example lookup table could be:

(174) TABLE-US-00003
TABLE 1 Graph Device Internet Protocol MAC Name Identifier (IP) Address
Address Cloud1 Cloud Server 84.15.186.115 00-17-FC-3F-FF-FF Cloud2 Remote Computing
84.15.184.60 00-13-DD-F5-AB-F5 System Edge1 Building Controller 192.168.1.16 00-16-23-FF-44-25
Edge2 AHU Controller 192.168.1.16 00-16-FC-22-56-21 Edge3 Data Collector 192.0.2.1 00-13-CA-44-
21-89 Edge4 Local Server 192.0.2.1 00-19-FC-39-A6-99

(175) Furthermore, each of the twin manager **108** and the edge device **1204** can store and execute agents, e.g., the twin manager **108** includes cloud agents **1206** while the edge device **1204** includes edge agents **1208**. The agents **1206** and **1208** can be the same as, or similar to, the agents described in FIGS. 5-9. By running agents on various cloud and/or edge devices and/or by storing the parts of a building graph across the various cloud and/or edge devices can allow for a digital twin (or digital twins) to be distributed across various cloud and/or edge devices. By having agents running against parts of the graph stored on the various cloud and/or edge devices can allow for digital twin computing and/or storage to co-exist at the same location where the computations are being used, providing autonomous execution capabilities. For example, if a controller for an AHU stores a portion of a building graph describing the AHU and an agent stored on the controller provides algorithms for running the AHU, the AHU controller can run a digital twin for controlling the AHU through the agent and the building graph stored by the AHU. This localizes the storage and execution of a digital twin at edge devices.

(176) This architecture shown in FIG. 12 also allow for flexibility of where the digital twin computations reside (e.g., on various clouds and/or edge devices). The architecture allows digital twin computations to not be co-located with the a cloud platform, e.g., the cloud platform **106** and/or the twin manager **108**. Because the agents and graph can be stored locally with the edge device **1204**, this can result in faster computing and/or control decisions since the computations are performed at the edge instead of requiring the edge device **1204** to wait for communication over the network **104** with the twin manager **108**.

(177) In some embodiments, one or multiple edge devices could store an edge graph that is not utilized unless a connection to a cloud system is offline and/or the cloud system is not operating properly. In this regard, the edge graph could be in a dormant mode where the edge devices store and/or maintain the graph but do not execute any applications against the graph, e.g., run a digital twin at the edge. The dormant graph could provide redundancy with the graph stored in the cloud system (e.g., include at least some of the same nodes and edges). The edge devices can perform operations based on a digital twin running in the cloud system. However if the edge devices become disconnected from the cloud system or the cloud system goes offline, edge systems can run an edge digital twin based on the edge graphs stored by the edge systems, e.g., “wake up” the edge digital twin.

(178) Referring now to FIG. 13, the distributed building graph **1100** generated across the twin manager **108** and an edge device **1204** during an onboarding process is shown, according to an exemplary embodiment. The edge device **1204**, the twin manager **108**, and the building controller **1306** can construct the building graph **1100** in a distributed manner. The edge device **1204**, the twin manager **108**, and the building controller **1306** include onboarding tools for constructing the graph, e.g., onboarding tool **1304**, onboarding tool **1308**, and onboarding tool **1312**.

(179) The onboarding tools can build the building graph **1100** in a piecemeal manner from the various device that the onboarding tools reside. While the building graph is constructed, each edge system can identify which pieces of the graph should be stored locally, e.g., which pieces the edge system needs to execute and which pieces of the constructed graph should be pushed to the cloud. The edge systems can further identify which portions of the graph to keep locally based on privacy. In some cases, the portions of the building graph that are retained locally at the edge can be linked via one or more nodes or edges with other portions of the building graphs stored by the other edge systems and/or cloud systems.

(180) In some embodiments, the onboarding tool **1308** of the twin manager **108** is pushed data collected by the edge device **1204** and data collected by the building controller **1306**. The twin manager **108** (e.g., via the onboarding tool **1308**) can generate a graph based on all of the data received. The twin manager **108** can make decisions of what portions of the graph should be pushed to the edge device **1204** and the building controller **1306**. In some embodiments, each onboarding tool **1304**, **1308**, and **1312** locally builds a graph to be stored locally on each of the edge device **1204**, the twin manager **108**, and the building controller **1306**. Similarly, in some embodiments, each edge device can identify what agents the edge device needs to execute to operate the equipment managed by the edge device. The edge devices can retrieve the agent from the twin manager **108**, in some embodiments. In some embodiments, the twin manager **108** can analyze the graph of each edge device (or a complete graph stored by the twin manager **108**) and determine what agents to push to the edge devices.

(181) In some embodiments, the onboarding tools **1304**, **1308**, and **1312** can continue managing and updating the graphs stored on the edge device **1204**, the twin manager **108**, and the building controller **1306** respectively. As devices of the building change (which may be infrequent) the onboarding tools **1304**, **1308**, and **1312** can perform updates to the various graphs and maintain consistent linkages across the various graphs.

(182) In FIG. 13, the building controller manages an AHU controller **1322** and a VAV1-2 controller **1316**. The AHU controller **1322** communicates with various AHU sensors **1324** and AHU actuators **1326**. The VAV1-2 controller communicates with VAV sensors **1318** and VAV actuators **1320**. The onboarding tool **1304** can receive metadata from the AHU controller **1322** and the metadata from the VAV1-2 controller **1316** defining the AHU sensors **1324**, the AHU actuators **1326**, the VAV sensors **1318**, and the VAV actuators **1320**. The onboarding tool **1304** can analyze the metadata to construct the edge building graph **1104**.

(183) The onboarding tool **1304** can perform various data analysis techniques to generate the edge building graph **1104** from the metadata, e.g., the techniques described in U.S. patent Ser. No. 16/885,959 filed May 28, 2020, U.S. patent Ser. No. 16/885,968 filed May 28, 2020, U.S. patent Ser. No. 16/663,623

filed Oct. 25, 2019, and U.S. patent Ser. No. 16/722,439 filed Dec. 20, 2019, the entirety of each of which are incorporated by reference herein.

(184) Furthermore, the onboarding tool **1304** can build the edge building graph **1104** based on information received from the twin manager **108** and/or the building controller **1306**. For example, the onboarding tool **1312**, while using data to build an edge building graph **1310** for the building controller **1306**, can communicate the data to the twin manager **108** and/or the edge device **1204**. In some embodiments, the twin manager **108** can communicate metadata received from the building controller **1306** to the edge device **1204**.

(185) In some embodiments, the data received by the edge device **1204** from the twin manager **108** is data describing the VAV**1-1** controller **1314** that communicates with the building controller **1306**. The data can provide an indication of the VAV**1-1**, an indication that the floor **1** and the floor**2** are fed by AHU**1**, and an indication that the zone**1-2** is part of the floor**1**. The onboarding tool **1304** can use the data to generate the edges **1162-1168**.

(186) Referring now to FIG. **14**, a process **1400** of querying a building graph divided into multiple building graphs distributed across cloud and edge systems is shown, according to an exemplary embodiment. The process **1400** can be performed by the twin manager **108** and the edge device **1204**. The process **1400** can be performed by one or more edge systems and/or one or more cloud systems. Furthermore, any computing device or group of computing devices can perform the process **1400**.

(187) In step **1402**, a system can receive a query for information of the building graph **1100** where the building graph **1100** is divided into multiple building graphs (e.g., the cloud building graph **1102** and the edge building graph **1104**) distributed across multiple device (e.g., the twin manager **108** and/or the edge device **1204**). In some embodiments, the query is an agent query generated by an agent, e.g., the cloud agents **1206** or the edge agents **1208**. The agent queries can be generated by the agents responsive to the agent identifying that it requires information not stored in a graph that the agent has direct access to (e.g., the cloud agents **1206** have direct access to the cloud building graph **1102**). In some embodiments, the queries are generated by applications (the applications **110**) that require some piece of information. In some embodiments, the queries are received, handled, and answered by the interfaces **1224** and **1226**.

(188) In step **1404**, the system can query a first building graph of the building graphs for the information and identify that at least a portion of the information is stored in a second building graph. For example, the interface **1224** could query the cloud building graph **1102** for information. In some embodiments, the interface **1224** retrieves at least a portion of the information from the cloud building graph **1102** to be returned in response to the query. However, the interface **1224** may further identify, based on the edges **1144**, **1146**, **1154**, **1160**, **1162**, **1164**, **1166**, and/or **1168**, that the information (or a portion of the information) needed to respond to the query is stored in the edge building graph **1104**.

(189) For example, if the query is for zones controlled by a VAV, the interface **1224** can identify the zone**1-1** node **1116** and identify that the zone**1-1** **1116** is fed by the VAV**1-2** **1114**. Furthermore, the interface **1224** can identify that the floor **1** node **1118** has a zone**1-2** via edge **1154** and identify based on the edge **1154** being related to the edge**1** node **1126**, the information describing the zone**1-2** is in the edge building graph **1104**.

(190) In step **1406**, the system can communicate the query to at least one of the cloud system or the edge system that stores the second building graph. For example, the interface **1226** may communicate a query to the interface **1224** for the interface **1224** to generate a result based on. Similarly, the interface **1224** may communicate a query to the interface **1226** for the interface **1226** to generate a result based on.

(191) Continuing the previous example, responsive to identifying that the building graph **1104** potentially stores information that should be returned in the query for zones controlled by a VAV, in the step **1406** the system, e.g., the interface **1224** can communicate the query to the interface **1226**. The interface **1224** may identify, based on the device-graph lookup table **1220** that the edge device **1204** stores the “Edge**1**” graph (e.g., the edge building graph **1104**), the “Edge**1**” graph indicated by the edge **1154**. The interface **1226** can query the edge building graph **1104** for zones controlled by VAVs.

(192) In step **1408**, the system can query the second building graph based on the query to identify the information and receive the information based on the query. Again, continuing the previous example, the interface **1226** can communicate the result of the query of the edge building graph **1104** back to the interface **1224**. The interface **1224** can compose a response (e.g., combine results from the graph **1102** and/or the graph **1104** or include results of only the graph **1102** and/or the graph **1104**). The response can

be returned by the interface **1224** to the requesting system, e.g., the cloud agents **1206**, applications, a user device, etc.

(193) Referring now to FIG. **15**, a process **1500** of querying a building graph divided into multiple building graphs distributed across cloud and edge systems by an agent, and ingesting information back into the building graph by the agent is shown, according to an exemplary embodiment. The process **1500** can be performed by the twin manager **108** and the edge device **1204**. The process **1500** can be performed by one or more edge systems and/or one or more cloud systems. Furthermore, any computing device or group of computing devices can perform the process **1500**.

(194) In step **1502**, an agent of the edge agents **1208** of the edge device **1204** can query the building graph **1100** divided into multiple building graphs (e.g., the cloud building graph **1102** and the edge building graph **1104**) and distributed across the cloud system (e.g., the twin manager **108**) and the edge device (e.g., the edge device **1204**) for information. The information can be information that the agent requires to perform operations, e.g., infer information, generate control decisions, run a temperature control algorithm, run an air quality algorithm, etc. In some embodiments, the agent queries the graphs **1102** and **1104** by providing the query to the interface **1226**. The interface **1226** can run the query, generate the query response, and provide the query response back to the agent.

(195) In step **1504**, the agent can receive first information from the first building graph stored by the cloud system and second information from a second building graph stored by the edge system. In some embodiments, the interface **1226** can compose a response including the first information and the second information and return the response to the agent. The cross-graph query can be the query described in the process **1400** of FIG. **14**.

(196) In step **1506**, the agent performs one or more operations to derive information based on the first information and the second information received in the step **1504**. In some embodiments, the agent may perform artificial intelligence (AI) based analysis to generate an event stream. The event stream may indicate a schedule of AI settings indicating the times at which settings should change. In step **1508**, the agent can ingest the information derived in the step **1506** into the first building graph or the second building graph. For example, the agent could ingest the AI inference event stream into the node **1134**. In some embodiments, the device manager **1218** and/or the agent can push the settings to the AHU (e.g., an actuator of the AHU such as the actuator device **1216**) to operate the AHU to control environmental conditions of a building (e.g., temperature, humidity, air quality, etc.).

(197) Referring now to FIG. **16**, a process **1600** is shown of querying a building graph divided into multiple graphs distributed across a cloud system and an edge platform and using information queried from the building graph to perform one or more operations, according to an exemplary embodiment. The process **1600** can be performed by the twin manager **108** and the edge device **1204**. The process **1600** can be performed by one or more edge systems and/or one or more cloud systems. Furthermore, any computing device or group of computing devices can perform the process **1600**.

(198) In step **1602**, a system (e.g., the twin manager **108** and/or the edge device **1204**) can query the building graph **1100** for information, the building graph **1100** divided into multiple building graphs distributed across a cloud system and an edge system, wherein one or more nodes or edges of a first graph stored by the edge system link the first building graph to a second building graph stored by another edge system or the cloud system. In some embodiments, the system performing the query can be an edge system, the cloud system, etc. In some embodiments, the query can be made across multiple graphs based on the edges that connect the graphs, e.g., the edges **1144**, **1146**, **1154**, **1160**, and/or **1162-1168**. In some embodiments, the system can identify which graph to query information from based on how the edges link the various graphs. The query can be performed across the building graphs as described in the process **1400** by FIG. **14**.

(199) In step **1604**, the system can receive the information from the building graph of the cloud system and/or the edge system. For example, the edge device **1204** could receive information of the edge building graph **1104** of the edge device **1204** or from the cloud building graph **1102** of the twin manager **108** or an edge building graph of another edge device. Similarly, the twin manager **108** could receive information from the cloud building graph **1102** and/or the edge building graph **1104** of the edge device **1204**.

(200) In step **1606**, the system can perform one or more operations on the information received in the step **1604** to derive an output and ingest the output into the building graph (e.g., the building graph **1102** or the edge building graph **1104**). The derived information can be displayed in the user device **176** for review by

a user, and/or use the information to operate building systems of the building, e.g., the building subsystems **122**, to control environmental conditions of a building, e.g., temperature, pressure, humidity, air quality, light, etc..

(201) Referring now to FIG. **17**, a process **1700** where the distributed building graph **1100** is generated across the twin manager **108** and an edge device **1204** during onboarding is shown, according to an exemplary embodiment. The process **1700** can be performed by the edge device **1204** and the twin manager **108**. The process **1700** can be performed by one or more edge systems and/or one or more cloud systems. Furthermore, any computing device described herein can be configured to perform the process **1700**.

(202) In step **1702**, the edge device **1204** can receive metadata from one or more devices connected to the edge device **1204**. For example, the metadata may be self-describing information of the one or more devices, e.g., device name, device purpose, device network identifier, device space location, sensor points, actuator points, units, data produced, etc. In some embodiments, the edge device **1204** receives, via the onboarding tool **1304**, metadata from the AHU controller **1322** describing the AHU controller **1322** and/or the AHU sensors **1324** and/or the AHU actuators **1326** that are connected to (and/or controlled by) the AHU controller **1322**. Similarly, in some embodiments, the edge device **1204** receives, via the onboarding tool **1304**, metadata from the VAV1-2 controller **1316** describing the VAV1-2 controller **1316** and/or the VAV sensors **1318** and/or the VAV actuators **1320** that are connected to (and/or controlled by) the VAV1-2 controller **1316**.

(203) In step **1704**, the edge device **1204** (e.g., via the onboarding tool **1304**) generates a first building graph, e.g., the edge building graph **1104**, which includes multiple nodes representing entities of the building including the one or more devices for which metadata is received in the step **1702**. The graph can further include edges that indicates the relationships between the entities. For example, based on the metadata received for the AHU controller **1322** (e.g., an identifier “AHU1”) and the metadata received for the VAV1-2 controller **1316** (e.g., an identifier “VAV1-2”), the onboarding tool **1304** could generate the respective the AHU1 node **1128** representing the AHU controlled by the AHU controller **1322** and the VAV1-2 node **1130** representing the VAV controlled by the VAV1-2 controller **1316**.

(204) Furthermore, the metadata received from the AHU controller **1322** and/or the VAV1-2 controller **1316** could indicate that the AHU1 feeds the VAV1-2 air. This could be identified through an explicit indication that the VAV1-2 receives air from the AHU1, through an identification of parameters controlled by the AHU controller affecting measurements of the VAV1-2, indications of network connections between the AHU controller **1322** and/or VAV1-2 controller **1316** indicating a control relationship, etc. The indication of the relationship can be used by the onboarding tool **1304** to generate the feeds edge **1170** linking the AHU1 node **1128** and the VAV1-2 node **1130**. The onboarding tool **1304** could further receive a location tag from the VAV1-2 controller **1316**. The location tag could indicate “zone1-2” or any other location name, e.g., “Main Lobby,” “Bob's Office,” “Family Room,” etc. The location tag may, in some embodiments, be programmed by a user for the VAV1-2 controller **1316** when the VAV1-2 controller **1316** is first installed. The onboarding tool **1304** can generate the zone1-2 node **1132** for the location of the VAV1-2 and generate a feeds edge **1172** between the VAV1-2 node **1130** and the zone1-2 node **1132** based on the indication that the VAV1-2 is installed in the zone1-2.

(205) In some embodiments, the edge agents **1208** may execute operations against the edge building graph **1104**. The results of the edge agents **1208** can, in some embodiments, be used by the edge agents **1208** and/or the onboarding tool **1304** to generate a node for the inferences produced by the edge agents **1208**. For example, if the edge agents **1208** derive information for the AHU1, the edge device **1204** can cause the edge building graph **1104** to include an AI inference event stream node **1134** and an edge **1174** between the AHU1 node **1128** and the node **1134**.

(206) In step **1706**, the edge device **1204** can receive, from a cloud system or another edge system, indications of second entities and second relationships between the entities and the second entities. The cloud system or another edge system may store a second building graph storing a plurality of second nodes representing the second entities and second edges between the second plurality of nodes representing relationships between the second entities. In some embodiments, the second building graph is generated in a similar manner as described in steps **1702-1704** by the cloud system or the other edge system. In step **1708**, the edge device **1204** can generate one or more nodes or edges linking the first building graph to the second building graph, the one or more nodes or edges relating the entities and the second entities.

(207) In some embodiments, the edge device **1204** the onboarding tool **1304** of the edge device **1204** may receive an indication of, or data of, the cloud building graph **1102** from the twin manager **108**. The onboarding tool **1304** can use the received data from the twin manager **108** to generate the node **1126** and/or the edges **1162-1168**. For example, in some embodiments, a name or other indicator of the cloud building graph **1102** received from the twin manager **108** can be used by the edge device **1204** to generate the cloud1 node **1126** representing the cloud building graph **1102**. Furthermore, the data received by the edge device **1204** from the twin manager **108** can indicate relationships between the nodes stored by the edge building graph **1104** and the nodes stored by the cloud building graph **1102**.

(208) For example, the edge device **1204** may receive an indication of a “VAV1-1” represented as a node in the cloud building graph **1102**. Furthermore, the edge device **1204** can receive an indication that the “VAV1-1” is fed by the AHU1 which is represented in the cloud building graph **1102**. Responsive to receiving this data, the edge device **1204** can generate the edge **1162** between the AHU1 node **1128** and the cloud1 node **1126**. Furthermore, the edge device **1204** could receive indications that two floors, floor **1** and floor **2** are also fed by the AHU1. Responsive to receiving this data, the edge device **1204** can generate the edges **1166** and **1162** between the AHU1 node **1128** and the cloud1 node **1126**. In some embodiments, the edge device **1204** can receive an indication that the cloud building graph **1102** includes a node that describes a floor that the zone1-2 represented by the zone1-2 node **1132** is a part of. Responsive to receiving this data, the edge device **1204** can generate the edge **1168** between the cloud1 node **1126** and the zone1-2 node **1132**.

(209) In some embodiments, a single system or device (e.g., the twin manager **108** and/or the edge device **1204**) could receive metadata from a variety of data sources, e.g., directly or indirectly from controllers, sensors, actuators, from one or more databases, from a user device, etc. The system can identify one or more edge systems or cloud systems that include processing resources (e.g., memory availability, storage availability, processing speed, number of processor cores, etc.) necessary to store and run a digital twin (e.g., a building graph and/or one or more agents). The system can construct multiple building graphs with interconnections. The system can communicate, via the network **104**, the building graphs to the one or more edge systems and/or cloud systems. In some embodiments, the twin manager **108** can construct the various graphs and push the graphs to the various edge devices. In some embodiments, the twin manager **108** may retain nodes or edges in the cloud that are linked to the nodes and/or edges of the graphs of the various edge devices.

(210) In some embodiments, a user may construct a digital twin via the user device **176**. In some embodiments, the user may provide input to the twin manager **108**, the user input defining one or multiple different graphs via the user device **176**. The user input can, in some embodiments, define the nodes and/or edges of each graph and interconnections between the graphs. The user input may further provide an assignment of each graph to a particular system, e.g., edge system, cloud system, etc. The twin manager **108**, responsive to receiving a user command, push the graphs to the various edge systems and/or cloud systems. Examples of constructing a graph via user input is described in U.S. patent application Ser. No. 16/723,087, filed Dec. 20, 2019 and U.S. patent application Ser. No. 16/175,507 filed Oct. 30 2018, the entireties of which are incorporated by reference herein.

CONFIGURATION OF EXEMPLARY EMBODIMENTS

(211) The construction and arrangement of the systems and methods as shown in the various exemplary embodiments are illustrative only. Although only a few embodiments have been described in detail in this disclosure, many modifications are possible (e.g., variations in sizes, dimensions, structures, shapes and proportions of the various elements, values of parameters, mounting arrangements, use of materials, colors, orientations, etc.). For example, the position of elements may be reversed or otherwise varied and the nature or number of discrete elements or positions may be altered or varied. Accordingly, all such modifications are intended to be included within the scope of the present disclosure. The order or sequence of any process or method steps may be varied or re-sequenced according to alternative embodiments. Other substitutions, modifications, changes, and omissions may be made in the design, operating conditions and arrangement of the exemplary embodiments without departing from the scope of the present disclosure.

(212) The present disclosure contemplates methods, systems and program products on any machine-readable media for accomplishing various operations. The embodiments of the present disclosure may be implemented using existing computer processors, or by a special purpose computer processor for an

appropriate system, incorporated for this or another purpose, or by a hardwired system. Embodiments within the scope of the present disclosure include program products comprising machine-readable media for carrying or having machine-executable instructions or data structures stored thereon. Such machine-readable media can be any available media that can be accessed by a general purpose or special purpose computer or other machine with a processor. By way of example, such machine-readable media can comprise RAM, ROM, EPROM, EEPROM, CD-ROM or other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to carry or store desired program code in the form of machine-executable instructions or data structures and which can be accessed by a general purpose or special purpose computer or other machine with a processor. When information is transferred or provided over a network or another communications connection (either hardwired, wireless, or a combination of hardwired or wireless) to a machine, the machine properly views the connection as a machine-readable medium. Thus, any such connection is properly termed a machine-readable medium. Combinations of the above are also included within the scope of machine-readable media. Machine-executable instructions include, for example, instructions and data which cause a general purpose computer, special purpose computer, or special purpose processing machines to perform a certain function or group of functions.

(213) Although the figures show a specific order of method steps, the order of the steps may differ from what is depicted. Also two or more steps may be performed concurrently or with partial concurrence. Such variation will depend on the software and hardware systems chosen and on designer choice. All such variations are within the scope of the disclosure. Likewise, software implementations could be accomplished with standard programming techniques with rule based logic and other logic to accomplish the various connection steps, processing steps, comparison steps and decision steps.

(214) In various implementations, the steps and operations described herein may be performed on one processor or in a combination of two or more processors. For example, in some implementations, the various operations could be performed in a central server or set of central servers configured to receive data from one or more devices (e.g., edge computing devices/controllers) and perform the operations. In some implementations, the operations may be performed by one or more local controllers or computing devices (e.g., edge devices), such as controllers dedicated to and/or located within a particular building or portion of a building. In some implementations, the operations may be performed by a combination of one or more central or offsite computing devices/servers and one or more local controllers/computing devices. All such implementations are contemplated within the scope of the present disclosure. Further, unless otherwise indicated, when the present disclosure refers to one or more computer-readable storage media and/or one or more controllers, such computer-readable storage media and/or one or more controllers may be implemented as one or more central servers, one or more local controllers or computing devices (e.g., edge devices), any combination thereof, or any other combination of storage media and/or controllers regardless of the location of such devices.

Claims

1. A method, comprising: retrieving, over a network by one or more processing circuits, metadata from one or more edge devices located in a building to generate a first digital twin of the building to store on a first memory device, the metadata identifying a first space, point, or piece of equipment of the building related to a second space, point, or piece of equipment of the building; generating, by the one or more processing circuits, the first digital twin of the building using the metadata and storing the first digital twin on the first memory device, the first digital twin comprising: a plurality of entities of the building and a plurality of relationships between the plurality of entities; receiving, from a second memory device by the one or more processing circuits, an identifier of a second digital twin stored by the second memory device and an identifier of an entity stored within the second digital twin representing the second space, point, or piece of equipment; causing, by the one or more processing circuits, the first digital twin to include: an entity including the identifier of the second digital twin; a relationship between an entity of the plurality of first entities stored within the first digital twin representing the first space, point, or piece of equipment and the entity of the first digital twin comprising the identifier of the second digital twin, the relationship comprising the identifier of the entity stored within the second digital twin to relate the entity representing

the first space, point, or piece of equipment and the entity representing the second space, point, or piece of equipment; querying, by the one or more processing circuits, the first digital twin based on query data; and performing, by the one or more processing circuits, one or more operations based on a query result of the querying.

2. The method of claim 1, comprising: identifying, by the one or more processing circuits, a particular system that stores the second digital twin based on a lookup table, wherein the lookup table stores a plurality of indications of digital twins including the first digital twin and the second digital twin and links between the digital twins and systems that store the digital twins.

3. The method of claim 1, comprising: identifying, by a processing circuit of a first system, a communication issue between the first system storing the first digital twin and a second system storing the second digital twin; and performing, by the processing circuit, one or more operations based on data determined based on the first digital twin.

4. The method of claim 1, comprising: executing, by the one or more processing circuits, an agent to: perform one or more operations based on the first digital twin; and ingest a result of the one or more operations into the first digital twin.

5. The method of claim 1, wherein: the second digital twin includes a second relationship that forms a connection between the second digital twin and the first digital twin.

6. The method of claim 1, wherein: the first digital twin includes a first graph structure including a plurality of first nodes representing a plurality of first spaces, points, or pieces of equipment and a plurality of first edges between the plurality of first nodes representing a plurality of first relationships between the plurality of first spaces, points, or pieces of equipment; and the second digital twin includes a second graph structure including a plurality of second nodes representing a plurality of second spaces, points, or pieces of equipment and a plurality of second edges between the plurality of second nodes representing a plurality of second relationships between the plurality of second spaces, points, or pieces of equipment.

7. The method of claim 6, wherein: the first graph structure includes a node identifying the second graph structure; and the first graph structure includes an edge between a first node of the plurality of first nodes representing the first space, point, or piece of equipment and the node identifying the second graph structure.

8. The method of claim 7, wherein the edge includes data including: a first indication of the entity representing the first space, point, or piece of equipment of the building; a second indication that the entity representing the first space, point, or piece of equipment is stored by the first graph structure; a third indication of the entity representing the second space, point, or piece of equipment; a fourth indication that the entity representing the second space, point, or piece of equipment is stored by the second graph structure; and a fifth indication of a relationship type of a plurality of relationship types describing a relationship between the entity representing the first space, point, or piece of equipment and the entity representing the second space, point, or piece of equipment.

9. The method of claim 1, comprising: querying, by a processing circuit of a first system, the first digital twin based on the query data, the one or more processing circuits comprising the processing circuit of the first system; identifying, by the processing circuit of the first system, that the query result is stored in the second digital twin based on the relationship; communicating, by the processing circuit of the first system, the query data to a second system; receiving, by the processing circuit of the first system, the query result from the second system; and performing, by the processing circuit of the first system, one or more operations based on the query result.

10. The method of claim 9, wherein: the first system comprises a cloud system stored off premises from the building; and the second system comprises an edge system stored on a premises of the building.

11. The method of claim 9, comprising: performing, by the one or more processing circuits, an onboarding to generate the first digital twin and the second digital twin by: receiving, by the processing circuit of the first system, first building data; generating, by the processing circuit of the first system, the first digital twin based on the first building data; receiving, by a second processing circuit of the second system, second building data; generating, by the second processing circuit of the second system, the second digital twin based on the second building data; identifying, by the processing circuit of the first system, the relationship by communicating with the second system; and causing, by the processing circuit of the first system, the first digital twin to include the relationship.

12. The method of claim 9, comprising: receiving, by a second processing circuit of the second system, the

query data from the processing circuit of the first system; querying, by the second processing circuit of the second system, the second digital twin with the query data to generate the query result; and transmitting, by the second processing circuit of the first system, the query result to the processing circuit of the first system.

13. A building system, comprising: one or more memory devices storing instructions thereon that, when executed by one or more processors, cause the one or more processors to: retrieve, over a network, metadata from one or more edge devices located in a building to generate a first digital twin of the building to store on a first memory device, the metadata identifying a first space, point, or piece of equipment of the building related to a second space, point, or piece of equipment of the building; generate the first digital twin of the building using the metadata and store the first digital twin on the first memory device, the first digital twin comprising: a plurality of entities of the building and a plurality of relationships between the plurality of entities; receive, from a second memory device, an identifier of a second digital twin stored by the second memory device and an identifier of an entity stored within the second digital twin representing the second space, point, or piece of equipment; cause the first digital twin to include: an entity including the identifier of the second digital twin; a relationship between an entity of the plurality of entities stored within the first digital twin and representing the first space, point, or piece of equipment and the entity of the first digital twin comprising the identifier of the second digital twin, the relationship comprising the identifier of the entity stored within the second digital twin to relate the entity representing the first space, point, or piece of equipment and the entity representing the second space, point, or piece of equipment; query the first digital twin based on query data; and perform one or more operations based on a query result of the querying.

14. The building system of claim 13, wherein: the first digital twin includes a first graph structure including a plurality of first nodes representing a plurality of first spaces, points, or pieces of equipment and a plurality of first edges between the plurality of first nodes representing a plurality of first relationships between the plurality of first spaces, points, or pieces of equipment; and the second digital twin includes a second graph structure including a plurality of second nodes representing a plurality of second spaces, points, or pieces of equipment and a plurality of second edges between the plurality of second nodes representing a plurality of second relationships between the plurality of second spaces, points, or pieces of equipment.

15. The building system of claim 14, wherein: the first graph structure includes a node identifying the second graph structure; and the first graph structure includes an edge between a first node of the plurality of first nodes representing the first space, point, or piece of equipment and the node identifying the second graph structure.

16. The building system of claim 15, wherein the edge includes data including: a first indication of the entity representing the first space, point, or piece of equipment of the building; a second indication that the entity representing the first space, point, or piece of equipment is stored by the first graph structure; a third indication of the entity representing the second space, point, or piece of equipment; a fourth indication that the entity representing the second space, point, or piece of equipment is stored by the second graph structure; and a fifth indication of a relationship type of a plurality of relationship types describing a relationship between the entity representing the first space, point, or piece of equipment and the entity representing the second space, point, or piece of equipment.

17. One or more non-transitory storage medium storing instructions thereon, that, when executed by one or more processors, cause the one or more processors to: retrieve, over a network, metadata from one or more edge devices located in a building to generate a first digital twin of the building to store on a first memory device, the metadata identifying a first space, point, or piece of equipment of the building related to a second space, point, or piece of equipment of the building; generate the first digital twin of the building using the metadata and storing the first digital twin on the first memory device, the first digital twin comprising: a plurality of entities of the building and a plurality of relationships between the plurality of entities; receive, from a second memory device, an identifier of a second digital twin stored by the second memory device and an identifier of an entity stored within the second digital twin representing the second space, point, or piece of equipment; cause the first digital twin to include: an entity including the identifier of the second digital twin; a relationship between an entity of the plurality of entities stored within the first digital twin representing the first space, point, or piece of equipment and the entity of the first digital twin comprising the identifier of the second digital twin, the relationship comprising the identifier of the entity

stored within the second digital twin to relate the entity representing the first space, point, or piece of equipment and the entity representing the second space, point, or piece of equipment; query the first digital twin based on query data; and perform one or more operations based on a query result of the querying.

18. The one or more non-transitory storage medium of claim 17, wherein the instructions cause the one or more processors to: perform an onboarding to generate the first digital twin and the second digital twin by: receiving first building data; generating the first digital twin based on the first building data; receiving second building data; generating the second digital twin based on the second building data; identifying the relationship; and causing the first digital twin to include the relationship.

19. The one or more non-transitory storage medium of claim 17, wherein: the first digital twin includes a first graph structure including a plurality of first nodes representing a plurality of first spaces, points, or pieces of equipment and a plurality of first edges between the plurality of first nodes representing a plurality of first relationships between the plurality of first spaces, points, or pieces of equipment; and the second digital twin includes a second graph structure including a plurality of second nodes representing a plurality of second spaces, points, or pieces of equipment and a plurality of second edges between the plurality of second nodes representing a plurality of second relationships between the plurality of second spaces, points, or pieces of equipment.
