

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386848
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Thomas; Frank

Method and system for persisting data

Abstract

A method and system for persisting data are provided. Batch data is periodically extracted via a computer system from at least one primary data source. Batch data is transformed via the computer system. The batch data is loaded, via the computer system, into a database. Real-time data is continuously received from at least one secondary data source. The received real-time data is indexed in an index. Business intelligence reports are generated from both the real-time data in the index and the batch data in the database.

Inventors:	Thomas; Frank (Toronto, CA)
Applicant:	Kinaxis Inc. (Ottawa, CA)
Family ID:	59851002
Assignee:	Kinaxis Inc. (Ottawa, CA)
Appl. No.:	18/521417
Filed:	November 28, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240095256 A1	Mar. 21, 2024

Related U.S. Application Data

continuation parent-doc US 17697317 20220317 US 11868363 child-doc US 18521417
continuation parent-doc US 16088149 US 11308115 20220419 WO PCT/CA2017/050330
20170314 child-doc US 17697317
us-provisional-application US 62307889 20160314

Publication Classification

Int. Cl.: **G06F16/25** (20190101); **G06F16/22** (20190101); **G06Q10/06** (20230101); **G06Q10/067** (20230101); **G06Q10/10** (20230101)

U.S. Cl.:

CPC **G06F16/254** (20190101); **G06F16/2228** (20190101); **G06F16/25** (20190101); **G06Q10/06** (20130101); **G06Q10/067** (20130101); **G06Q10/10** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7461039	12/2007	Gilman	706/25	G06F 16/254
8930380	12/2014	Saurabh et al.	N/A	N/A
8978034	12/2014	Goodson et al.	N/A	N/A
9665621	12/2016	Reiner et al.	N/A	N/A
10424020	12/2018	Kmak et al.	N/A	N/A
10452981	12/2018	Christie, IV et al.	N/A	N/A
2002/0174126	12/2001	Britton	N/A	G06Q 30/02
2003/0084053	12/2002	Govrin et al.	N/A	N/A
2004/0083199	12/2003	Govindugari et al.	N/A	N/A
2005/0251812	12/2004	Hayward	N/A	N/A
2005/0262192	12/2004	Mamou et al.	N/A	N/A
2005/0289138	12/2004	Cheng et al.	N/A	N/A
2007/0160067	12/2006	Hayward	N/A	N/A
2008/0046536	12/2007	Broda	N/A	N/A
2011/0035744	12/2010	Bhatia	N/A	N/A
2011/0301982	12/2010	Green, Jr. et al.	N/A	N/A
2011/0313969	12/2010	Ramu	N/A	N/A
2012/0102053	12/2011	Barrett et al.	N/A	N/A
2012/0143715	12/2011	Eshghi et al.	N/A	N/A
2013/0185349	12/2012	Broda	N/A	N/A
2014/0095645	12/2013	Grosz	715/753	G06F 3/1272
2014/0229628	12/2013	Mandal	709/231	G06F 16/24568
2014/0289188	12/2013	Shimanovsky et al.	N/A	N/A
2015/0205602	12/2014	Prismon et al.	N/A	N/A
2016/0055499	12/2015	Hawkins et al.	N/A	N/A
2016/0162550	12/2015	Wang et al.	N/A	N/A
2016/0234087	12/2015	Nyerges	N/A	H04L 45/02
2017/0006135	12/2016	Siebel et al.	N/A	N/A
2017/0046409	12/2016	Bender	N/A	N/A
2017/0076323	12/2016	Palanciuc et al.	N/A	N/A
2018/0084073	12/2017	Walsh et al.	N/A	N/A

OTHER PUBLICATIONS

Bornea et al., "Semi-Streamed Index Join for Near-Real Time Execution of ETL Transformations," Proceedings of the 2011 IEEE 27th International Conference on Data Engineering ICDE 2011, Apr. 2011, pp. 159-170. cited by applicant

European Search Report for 17765609.7, EPO, dated Sep. 6, 2019. cited by applicant

International Patent Application No. PCT/CA2017/050330, International Preliminary Report on Patentability dated Sep. 27, 2018. cited by applicant

International Search Report corresponding to PCT/CA2017/050330; Canadian Intellectual Property Office; issued Jul. 21, 2017. cited by applicant

Santos et al., "Real-Time Data Warehouse Loading Methodology," Proceedings of the 2008 International Symposium on Database Engineering Applications IDEAS '08, Sep. 2008, pp. 49-58. cited by applicant

U.S. Appl. No. 16/088,149, Notice of Allowance dated Dec. 22, 2021. cited by applicant

U.S. Appl. No. 17/697,317, Advisory Action dated Jul. 25, 2023. cited by applicant

U.S. Appl. No. 17/697,317, Final office action dated May 2, 2023. cited by applicant

U.S. Appl. No. 17/697,317, Non Final office Action dated Oct. 12, 2023. cited by applicant

U.S. Appl. No. 17/697,317, Non Final office Action dated Oct. 18, 2022. cited by applicant

U.S. Appl. No. 17/697,317, Notice of Allowance dated Nov. 9, 2023. cited by applicant

U.S. Appl. No. 16/088,149, Final Office Action dated Feb. 9, 2021. cited by applicant

U.S. Appl. No. 16/088,149, Non-Final Office Action dated Jul. 23, 2021. cited by applicant

U.S. Appl. No. 16/088,149, Non-Final Office Action dated Jun. 19, 2020. cited by applicant

U.S. Appl. No. 16/088,149, Non-Final Office Action dated Nov. 6, 2020. cited by applicant

Written Opinion of the International Searching Authority corresponding to PCT/CA2017/050330; Canadian Intellectual Property Office; issued Jul. 21, 2017. cited by applicant

Primary Examiner: Bowen; Richard L

Attorney, Agent or Firm: Borden Ladner Gervais LLP

Background/Summary

(1) This application is a continuation of U.S. patent application Ser. No. 17/697,317, filed on Mar. 17, 2022, which is a continuation of U.S. patent application Ser. No. 16/088,149, filed Sep. 25, 2018, now issued as U.S. Pat. No. 11,308,115 on Apr. 19, 2022, which is a 371 of international PCT/CA2017/050330, filed Mar. 14, 2017, which claims benefit of U.S. provisional patent application Ser. No. 62/307,889, filed Mar. 14, 2016, the contents of which are incorporated entirely herein by reference.

TECHNICAL FIELD

(1) The following relates generally to informatics and more specifically to a method and system for persisting data.

BACKGROUND

(2) Business intelligence ("BI") systems are known. Such systems are used to acquire and transform raw data into meaningful and useful information for business analysis purposes. They are capable of handling large amounts of unstructured data to help identify, develop, and otherwise create new business opportunities. Such systems allow the data to be visualized in ways that provide insight into trends and other characteristics of the data.

(3) Prior to analysis, the data is batch extracted from one or more data sources by an extract-

transform-load (“ETL”) engine and stored in a database, such as a flat file or a data warehouse. During the ETL process, data is extracted from the one or more data sources such as operational databases. The data sources may be homogeneous or may use different data structures and/or formats. Data validation may be performed during the extraction phase. Next, the ETL engine transforms the extracted data using a set of rules or functions to prepare it for loading into the data warehouse. Once the extracted data is transformed, it can then load the data into the database. During the loading of the data, the data in the repository may be overwritten, appended to, updated, etc.

(4) As the ETL process can take a significant amount of time, and both the data sources and the data warehouse can suffer reduced performance, it is generally performed periodically, such as once a day, week, or month. As a result, the BI reports generated by the system are generally not current.

(5) Current real-time data is often collected in a separate database and analyzed separately. Analysis of the aggregate data stored in the separate database together with the batch data can be quite manual and, as a result, can be very time-consuming and error-prone.

SUMMARY

(6) In an aspect, a method for persisting data is provided, the method comprising periodically extracting batch data via a computer system from at least one primary data source, transforming, via the computer system, the batch data, loading, via the computer system, the batch data into a database, continuously receiving real-time data from at least one secondary data source, indexing the received real-time data in an index, and generating business intelligence reports from both the real-time data in the index and the batch data in the database.

(7) In a particular case, the method further comprising storing the extracted batch data in an interim data store.

(8) In another case, transforming the batch data comprises transforming the batch data into a format specified by a data schema.

(9) In yet another case, transforming the batch data comprises at least one of discarding certain data, translating the data to account for differences in a representation of the data, determining the age of the data, sorting the data, joining the data where the data is from multiple sources, and aggregating the data.

(10) In yet another case, the at least one primary data source does not publish or push the data immediately.

(11) In yet another case, the method further comprising analyzing the data in the database for generating the business intelligence reports, the analyzing comprising at least one of statistical analysis, machine learning techniques, and data mining.

(12) In yet another case, the method further comprising periodically reconciling the real-time data against the batch data.

(13) In yet another case, the real-time data is acquired via at least one of REST requests and a user datagram protocol.

(14) In yet another case, the real-time data is pushed onto a distributed message queue system.

(15) In yet another case, the method further comprising storing the real-time data by key-value in a distributed key-value store for subsequent integration into the database.

(16) In another aspect, there is provided a system for persisting data comprising one or more processors and a data storage, the one or more processors configured to execute, or direct to be executed: an extract-transform-load module to: periodically extract batch data from at least one primary data source; transform the batch data; and load the batch data into a database; a real-time distributed processing module to: continuously receive real-time data from at least one secondary data source; index the received real-time data to generate an index; and at least one business intelligence tool to generate business intelligence reports from both the real-time data in the index and the batch data in the database.

- (17) In a particular case, the extract-transform-load module further executable to store the extracted batch data in an interim data store.
- (18) In another case, transforming the batch data by the extract-transform-load module comprises transforming the batch data into a format specified by a data schema.
- (19) In yet another case, transforming the batch data by the extract-transform-load module comprises at least one of discarding certain data, translating the data to account for differences in a representation of the data, determining the age of the data, sorting the data, joining the data where the data is from multiple sources, and aggregating the data.
- (20) In yet another case, the at least one primary data source does not publish or push the data immediately.
- (21) In yet another case, the system further comprising a data analysis module executable to analyze the data in the database for generating the business intelligence reports, the analysis comprising at least one of statistical analysis, machine learning techniques, and data mining.
- (22) In yet another case, the at least one business intelligence tool periodically reconciles the real-time data against the batch data.
- (23) In yet another case, the real-time data is acquired by the extract-transform-load module via at least one of REST requests and a user datagram protocol.
- (24) In yet another case, the system further comprising a distributed message queue system, the extract-transform-load module pushes the real-time data onto the distributed message queue system.
- (25) In yet another case, the real-time distributed module stores the real-time data by key-value in a distributed key-value store for subsequent integration into the database.
- (26) These and other aspects are contemplated and described herein. It will be appreciated that the foregoing summary sets out representative aspects of systems and methods to assist skilled readers in understanding the following detailed description.
-

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The features of the invention will become more apparent in the following detailed description in which reference is made to the appended drawings wherein:
- (2) FIG. 1 is a schematic diagram of a system for persisting data in accordance with an embodiment;
- (3) FIG. 2A is a view of the data persisted by the system of FIG. 1 generated using both batch data and real-time data;
- (4) FIG. 2B is a view of the data persisted by the system of FIG. 1 similar to that of FIG. 2A, with the incorporation of further batch data; and
- (5) FIG. 3 is a view of the data persisted by the system of FIG. 1 similar to that of FIG. 2A, with the incorporation of further batch data; and

DETAILED DESCRIPTION

(6) Embodiments will now be described with reference to the figures. For simplicity and clarity of illustration, where considered appropriate, reference numerals may be repeated among the Figures to indicate corresponding or analogous elements. In addition, numerous specific details are set forth in order to provide a thorough understanding of the embodiments described herein. However, it will be understood by those of ordinary skill in the art that the embodiments described herein may be practiced without these specific details. In other instances, well-known methods, procedures and components have not been described in detail so as not to obscure the embodiments described herein. Also, the description is not to be considered as limiting the scope of the embodiments described herein.

(7) Various terms used throughout the present description may be read and understood as follows, unless the context indicates otherwise: “or” as used throughout is inclusive, as though written “and/or”; singular articles and pronouns as used throughout include their plural forms, and vice versa; similarly, gendered pronouns include their counterpart pronouns so that pronouns should not be understood as limiting anything described herein to use, implementation, performance, etc. by a single gender; “exemplary” should be understood as “illustrative” or “exemplifying” and not necessarily as “preferred” over other embodiments. Further definitions for terms may be set out herein; these may apply to prior and subsequent instances of those terms, as will be understood from a reading of the present description.

(8) Any module, unit, component, server, computer, terminal, engine or device exemplified herein that executes instructions may include or otherwise have access to computer readable media such as storage media, computer storage media, or data storage devices (removable and/or non-removable) such as, for example, magnetic disks, optical disks, or tape. Computer storage media may include volatile and non-volatile, removable and non-removable media implemented in any method or technology for storage of information, such as computer readable instructions, data structures, program modules, or other data. Examples of computer storage media include RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store the desired information and which can be accessed by an application, module, or both. Any such computer storage media may be part of the device or accessible or connectable thereto. Further, unless the context clearly indicates otherwise, any processor or controller set out herein may be implemented as a singular processor or as a plurality of processors. The plurality of processors may be arrayed or distributed, and any processing function referred to herein may be carried out by one or by a plurality of processors, even though a single processor may be exemplified. Any method, application or module herein described may be implemented using computer readable/executable instructions that may be stored or otherwise held by such computer readable media and executed by the one or more processors.

(9) The following relates generally to informatics and more specifically to a method and a system for persisting data. The method and system provides a method by which real-time data can be analyzed alongside data that is periodically extracted, transformed, and loaded in a batch. Real-time data is acquired as it is generated, typically via push.

(10) Referring now to FIG. 1, a system **20** for persisting data in accordance with an embodiment is shown. The system **20** obtains both batch and real-time data from various data sources and stores the data to facilitate the consolidated reporting and visualization of the data. The system **20** acts as an end-to-end solution that is capable of consuming arbitrary data from various sources, transform it to a standard (but flexible) data model, and provides a means to consume the data from a BI reporting perspective, as well as provide a feature set with which to train machine learning models. The feature set is a set of data points used by a model. For example, the feature set can be a list of numbers that form a time series.

(11) By combining real-time and batch data, the system **20** can leverage the advantages of each simultaneously. For example, real-time systems greatly reduce the time to insight or can be used in real-time predictions (e.g., if it is desired to make a recommendation as a user is browsing an e-commerce site). Batch systems are typically more correct than real-time systems. This is because real-time data may be lost, and there are no strong ordering guarantees for real-time data (e.g., events from a user browsing a website may not be processed in the same order as they occurred). By combining the storage and processing of both batch data and real-time data, analysts can be provided quicker access to fresh data and any data integrity problems can be resolved through the batch jobs.

(12) The illustrated system **20** stores e-commerce transaction data. Real-time information may be

coming in from a web beacon injected into webpages generated by an ecommerce store. When executed by a browser of a customer, the web beacon collects purchase information and submits it to the system **20** via REST calls. Transaction information is collected directly from the e-commerce platform, which is typically only extracted from the operational data stores, transformed, and loaded into a database, such as an analytical database, periodically.

(13) The real-time information from the web beacon can be used to create visualizations that display sales volume by region as the sales happen, as well as run models against the incoming data for fraud detection, generate popups displaying product recommendations, and tailor the website for the user based on browsing patterns. When the data is stored, there may be some discrepancies due to the nature of using real-time data and systems. The real-time data can be reconciled against the batch data extracted from the e-commerce site on a schedule (e.g., nightly).

(14) The system **20** includes an ETL module **24** that is in communication with a set of one or more data sources **28a** to **28d**. The data sources **28a** to **28d** (collectively **28**) can be commercial or open-source databases, text files and sources, and any other data source type from which data can be extracted. In the illustrated embodiment, the data sources include an Open Database Connectivity (“ODBC”)-compatible database **28a**, a NoSQL database **28b** (e.g., MongoDB, Cassandra, etc.) that is non-relational, a raw text file **28c** hosted in any location (such as a file of comma-separated values on a hard drive), and a custom data source **28d**, such as a web log or an analytics source. The data sources **28** are not real-time in nature, in that changes made to the data sources are not published or otherwise pushed immediately, and thus must be pulled from these data sources.

(15) The ETL module **24** pulls or extracts data from the data sources **28** in batches via a set of data extraction tools **32a** to **32c**. These batch extractions are automatically performed periodically, typically daily or weekly at night to minimize the impact on the availability of the data sources **28**.

(16) The batch data extracted by the data extraction tools **32a** to **32c** is dumped into an interim data store **36**. The batch data extracted from the data sources **28** is typically not reformatted in any way and, as a result, the format of the extracted data in the interim data store **36** is not markedly different than that of the data sources **28**.

(17) An ETL engine **40** then processes the data from each data source generally separately to transform the batch data into a format specified by a data schema for a database **44**. Transformation scripts that are used by the ETL engine **40** to transform the batch data include a set of rules are used to specify the transformations to be applied. During the transformation, the batch data may be cleaned to filter out unwanted data. Examples of transformations performed may include the following functions: discarding certain data items, translating coded data items to account for differences in the representation of data between the data source and the data warehouse (for example, gender may be represented in a data source as “F” or “M”, whereas it may be represented as “1” or “2” in the data warehouse), calculating data items such as age from the current date and a date of birth, sorting or ordering the data, joining data from multiple sources, aggregating data to generate a rollup, etc.

(18) Once the extracted batch data is transformed, the ETL engine **40** loads the batch data into the database **44**. The database **44** can be any one of a number of data store types that permits querying of the data stored therein, but in the illustrated embodiment is a massive parallel processing (“MPP”) relational database. Preferably, the database **44** can scale horizontally, and has a shared nothing architecture (e.g., Greenplum).

(19) The batch data from each data source **28** can be separately transformed and loaded into the database **44**, or may be collectively loaded into the database **44** once transformed. During the loading of the batch data, the batch data in the repository may be overwritten, appended to, updated, etc.

(20) The database **44** acts as the layer that makes the batch data accessible via queries.

(21) The system **20** also acquires real-time data. The real-time data can be acquired from any data sources that are real-time in nature. In the illustrated embodiment, the system **20** acquires real-time

data in the form of REST requests received via a RESTful web server **48**. The REST calls can be sent via a JavaScript beacon inserted into a webpage for a web application or ecommerce website. These requests could contain any arbitrary bit of information such as web events, purchases, users interacting with a website, search queries, etc. As will be understood, the REST calls represent data that is pushed into the system **20** from other systems.

(22) Data received via the RESTful web server **48** is pushed onto a distributed message queue system **52**. The distributed message queue system **52** provides a means for many concurrent workers to crunch through the data simultaneously. It also enables a layer of reliability—if a message fails in transit or times out, the distributed message queue system **52** detects this and replays the message.

(23) The messages in the distributed message queue system **52** are retrieved and processed via a real-time distributed processing framework **56** by a real-time distributed processing module **57**. The real-time distributed processing framework **56** is a mechanism that can process data in a distributed fashion. The real-time distributed processing framework **56** can scale horizontally by simply adding additional workers.

(24) The real-time distributed processing framework **56**, on the real-time distributed processing module **57**, processes data in the distributed message queue system **52** by providing the messages representing the real-time data to a search engine **60** and a distributed key-value store **64**. The search engine **60** provides an easy mechanism to index fast moving data as well as produce simple aggregations in real-time, and is used to power real-time BI tools.

(25) The distributed key-value store **60** stores the real-time data received for later integration into the database **44**. An exemplary distributed key-value store is Apache Cassandra. Real-time data is stored by the distributed key-value store **60** by key-value. The advantage of using a key-value store is that if a message is processed multiple times, the key-store can identify and handle duplicates accordingly. Time series information can easily be stored in a key-value store as a specific time range can be specified to bucket the data as a key (e.g., minute, hour), the timestamp as a column (to keep the data sorted). All of the messages that occurred during that time frame can be retrieved in order by simply querying for that bucket.

(26) Various data analysis modules **68** can process the data stored in the database **44** to summarize the data and draw inferences. The data analysis modules **68** can perform statistical analysis, machine learning techniques, data mining, or any other sort of data processing. The results generated by the data analysis modules **68** are then placed into the database **44** to enable the BI tools to access, process and/or present the results.

(27) A BI visualization layer **72** includes at least one of business intelligence tools **71** for business intelligence, reporting, and/or visualization. These tools **71** enable a person to visualize the raw and processed batch data stored in the database **44** and the real-time data stored by the search engine **60**. In this manner, the BI visualization layer **72** enables a person to visualize raw data, reports, visualization tools for both the batch data and the real-time data received by the system **20**, as well as inferences generated by the data analysis modules **68** for the batch data.

(28) The ETL engine **40** periodically (nightly in this embodiment) extracts data from the distributed key-value store **64** and places it in the interim data store **36**. Next, the ETL engine **40** transforms the data from the distributed key-value store in the interim data store **36** into a format specified by the data schema for the database **44**. Transformation scripts that are used by the ETL engine **40** to transform the batch data include a set of rules are used to specify the transformations to be applied to the real-time data. During the transformation, some real-time data items already available in the batch data can be discarded, as the batch data may be assumed to be more reliable. Other data items, such as user interactions with a website, are kept and aligned with the batch data in the database **44** via one or more common keys, such as purchase order number, etc. Then, the ETL engine **40** loads the transformed data into the database **44**. In this manner, the database **44** is furnished with batch data from the data sources **28** and supplemented with real-time data received

via the RESTful web server **48**.

(29) The system **20** uses batch data, which is generally more accurate/reliable than real-time data, where possible, and supplements the batch data with real-time data when batch data is unavailable.

(30) FIG. 2A shows a chart **100** representing a view of the data generated by a visualization tool **71** in the BI visualization layer **72** of the system **20** of FIG. 1. As illustrated, the data includes both batch data from the database **44** for the period of time up to the last batch extraction from the data sources **28** and real-time data from the index managed by the search engine **60** for the period after this point.

(31) FIG. 2B shows a chart **100'** presenting the same view of the data generated by the visualization tool **71** of the BI visualization layer **72** of the system **20** shown in FIG. 2A, but includes additional batch data after a subsequent batch data extraction, transformation, and loading into the database **44**. As can be seen, the batch data for the period since the last extraction to the newer extraction time replaces the real-time data for the same period, and real-time data from the index is used after this point.

(32) FIG. 3 shows a chart **100'** presenting the same view of the data generated by the visualization tool **71** of the BI visualization layer **72** of the system **20** shown in FIG. 2A, but includes additional batch data after a subsequent batch data extraction, transformation, and loading into the database **44**. As can be seen, the batch data for the period since the last extraction to the newer extraction time replaces the real-time data for the same period, and real-time data from the index is used after this point.

(33) The above approach reduces the deployment and maintenance burden, as there is one integrated system with a single management console rather than several for handling the different types of data. A standard interface for the application layer (data model) can be provided. Further, the approach is flexible, enabling users to rapidly build new use-cases for the system as well as iterate on existing ones. Flexibility is achieved through consolidating several technologies into one platform.

(34) While the system described above uses a RESTful web server to receive real-time data, those of skill in the art will appreciate that there are numerous other solutions for receiving real-time data. For example, some systems can be configured to receive pushed data via user datagram protocol (“UDP”).

(35) The pushed data can be received from a variety of sources. Internet-enabled appliances, such as thermostats, set top boxes, home automation hubs, etc., as well as more traditional computing devices, such as desktop computers, servers, mobile phones, tablets, etc. can be configured to push data to the system.

(36) In embodiments, the components of the system are stored by and executed on a single computer system or a single cluster of computers. In other embodiments, the components of the system **20** are distributed among two or more computer systems or clusters of computer systems that may be locally or globally distributed. The system provides a means for operating multiple types of database systems, data processing frameworks, operating systems, and hardware systems remotely. This management system oversees the computer system(s) or cluster(s) executing or storing the individual databases, search engine, ETL machines, etc. The management system makes calls to application programming interfaces (“APIs”) of these subsystems and/or modifies their configuration (either through files or via an API) in order to perform tasks such as backups, configuration modification, etc. In addition, the management system can use third-party monitoring tools and expose them through the console.

(37) Although the invention has been described with reference to certain specific embodiments, various modifications thereof will be apparent to those skilled in the art without departing from the spirit and scope of the invention as outlined in the claims appended hereto. The entire disclosures of all references recited above are incorporated herein by reference.

Claims

1. A method for persisting and reconciling one or more data sources on a computer system, the method comprising: periodically extracting batch data from at least one primary data source, transforming the batch data into a first format specified by a data schema for a database; loading the transformed batch data into the database continuously receiving real-time data from at least one secondary data source; indexing the received real-time data to generate an index; generating business intelligence reports comprising the real-time data in the index and the transformed batch data in the database; storing the real-time data by key-value in a distributed key-value store for subsequent integration into the database, the key-value bucketing a time range of the real-time data having time-series information within the time range of the key-value, and transforming the real-time data stored in the distributed key-value store into the first format specified by the data schema for the database.
2. The method of claim 1, further comprising storing the extracted batch data in an interim data store.
3. The method of claim 1, wherein transforming the batch data comprises at least one of discarding certain data, translating the data to account for differences in a representation of the data, determining the age of the data, sorting the data, joining the data where the data is from multiple sources, and aggregating the data.
4. The method of claim 1, wherein the at least one primary data source does not publish or push the data immediately.
5. The method of claim 1, further comprising analyzing the data in the database for generating the business intelligence reports, the analyzing comprising at least one of statistical analysis, machine learning techniques, and data mining.
6. The method of claim 1, further comprising periodically reconciling the real-time data against the batch data.
7. The method of claim 1, wherein the real-time data is acquired via at least one of REST requests and a user datagram protocol.
8. The method of claim 1, wherein the real-time data is pushed onto a distributed message queue system.
9. A system for persisting and reconciling one or more data sources, the system comprising one or more processors and a data storage, the one or more processors configured to execute, or direct to be executed: an extract-transform-load module to: periodically extract batch data from at least one primary data source; transform the batch data into a first format specified by a data schema for a database; and load the transformed batch data into the database; a real-time distributed processing module to: continuously receive real-time data from at least one secondary data source; index the received real-time data to generate an index; store the real-time data by key-value in a distributed key-value store for subsequent integration into the database, the key-value bucketing a time range of the real-time data having time-series information within the time range of the key-value; and transform the real-time data stored in the distributed key-value store into the first format specified by the data schema for the database, and at least one business intelligence tool to generate business intelligence reports comprising the real-time data in the index and the transformed batch data in the database.
10. The system of claim 9, the extract-transform-load module further executable to store the extracted batch data in an interim data store.
11. The system of claim 9, wherein transforming the batch data by the extract-transform-load module comprises at least one of discarding certain data, translating the data to account for differences in a representation of the data, determining the age of the data, sorting the data, joining the data where the data is from multiple sources, and aggregating the data.

12. The system of claim 9, wherein the at least one primary data source does not publish or push the data immediately.

13. The system of claim 9, further comprising a data analysis module executable to analyze the data in the database for generating the business intelligence reports, the analysis comprising at least one of statistical analysis, machine learning techniques, and data mining.

14. The system of claim 9, wherein the at least one business intelligence tool periodically reconciles the real-time data against the batch data.

15. The system of claim 9, wherein the real-time data is acquired by the extract-transform-load module via at least one of REST requests and a user datagram protocol.

16. The system of claim 9, further comprising a distributed message queue system, the extract-transform-load module pushes the real-time data onto the distributed message queue system.
