

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12386898
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Schilders; Steven

Systems and methods for sub-graph processing in executable graph-based models

Abstract

A method for dynamic execution of sub-graphs within executable graph-based models is provided. Processing circuitry obtains an executable graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs. Each sub-graph defines a hierarchical structure of related nodes. The processing circuitry receives a stimulus and a context associated with the stimulus. In response to the stimulus being received and based on the context, the processing circuitry maps the stimulus to a first sub-graph of the executable graph-based model. The processing circuitry causes execution of processing logic within the overlay structure based on the mapping. The processing logic is associated with one or more nodes of the first sub-graph.

Inventors:	Schilders; Steven (Columbus, IN)
Applicant:	INFOSYS LIMITED (Bangalore, IN)
Family ID:	85772134
Assignee:	Infosys Limited (Bangalore, IN); InvertIT Inc. (Columbus, IN)
Appl. No.:	18/125540
Filed:	March 23, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240256603 A1	Aug. 01, 2024

Related U.S. Application Data

us-provisional-application US 63448718 20230228
us-provisional-application US 63442682 20230201

Publication Classification

Int. Cl.: G06F16/90 (20190101); G06F16/901 (20190101)

U.S. Cl.:

CPC G06F16/9024 (20190101); G06F16/9027 (20190101);

Field of Classification Search

CPC: G06F (16/9024); G06F (16/9027)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11815943	12/2022	Rosendahl	N/A	G06F 9/547
12081636	12/2023	Enrici	N/A	H04L 67/60
2023/0023655	12/2022	Hui et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2018111699	12/2017	WO	N/A

OTHER PUBLICATIONS

Europe Patent Application No. 23165125.8, Extended European Search Report, dated Dec. 11, 2023. cited by applicant

Bulatov: "Fitting larger networks into memory"; dated Jan. 14, 2018, URL:<https://medium.com/tensorflow/fittinglarger-networks-into-memory-583e3c758ff9> [retrieved on Nov. 24, 2023]. cited by applicant

Primary Examiner: Andersen; Kristopher

Attorney, Agent or Firm: Finnegan, Henderson, Farabow, Garrett & Dunner LLP

Background/Summary

RELATED APPLICATIONS (1) The present application claims priority to U.S. Patent Application Ser. No. 63/442,682, filed Feb. 1, 2023, and U.S. Patent Application Ser. No. 63/448,718, filed Feb. 28, 2023, which are incorporated herein by reference in their entirety.

TECHNICAL FIELD

(1) The present disclosure relates to executable graph-based models. Particularly, but not exclusively, the present disclosure relates to the hierarchical decomposition of executable graph-based models into sub-graphs. Particularly, but not exclusively, the present disclosure relates to the loading, management, and execution of sub-graphs within executable graph-based models.

BACKGROUND

(2) Modern system designs typically separate data storage from any functional data structure used

from a processing logic perspective. This separation often occurs when data is “at rest” or at run-time where the processing system interacts with a copy of the relevant data in the processing space that can be of a different representation. This separation also leads to an impedance mismatch which requires some form of a data management solution to perform the necessary mappings between the two states. As a result of this separate of concerns, the processing logic is typically performed in a separate technology and physical tier (in an n-tier architecture) from the data. This is illustrated in the example n-tier architecture shown in FIG. 1.

(3) The example n-tier architecture **100** comprises a presentation layer **102**, a processing logic layer **104**, a data access layer **106**, and a database layer **108**. The presentation layer **102** comprises applications or components which are used to display the outputs of the processing logic layer **104** to a user or users. The processing logic layer **104** comprises applications, components, or services which perform some form of processing on the data obtained from the data access layer **106**. The data access layer **106** comprises the applications, components, and/or services which can access the data used by the processing logic layer **104** and stored at the database layer **108**. The database layer **108** handles the persistent storage of the data used by the system (e.g., in the form of a relational database, flat file, NoSQL database, graph database, and the like).

(4) The layers of the example n-tier architecture **100** are technically separated. Each layer may utilize a separate set of components to perform specific functionality (e.g., a database management system is used in the database layer **108** whilst an enterprise application is used in the processing logic layer **104**). The layers of the n-tier architecture **100** may also be physically separated. For example, the database layer **108** may execute on a remote cloud service, the processing logic layer **104** may execute on a network within an enterprise, and the presentation layer **102** may execute on a user device within the enterprise. While some architectural designs require a clear separation of concerns between data and the use of the data, often the separation enforced by architectures such as that illustrated in FIG. 1 can severely inhibit the flexibility, extensibility, and responsiveness of any system created.

(5) Therefore, there is a need for enhanced architectures which provide improved flexibility, extensibility, and responsiveness thereby providing more efficient data processing systems.

SUMMARY OF DISCLOSURE

(6) According to an aspect of the present disclosure, there is provided a method for dynamic execution of sub-graphs within executable graph-based models. The method comprises obtaining an executable graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein each sub-graph defines a hierarchical structure of related nodes. The method further comprises receiving a stimulus and a context associated with the stimulus, and in response to the stimulus being received, mapping, based on the context, the stimulus to a first sub-graph of the executable graph-based model. The method further comprises causing execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with one or more nodes of the first sub-graph.

(7) According to a further aspect the present disclosure, there is provided a device comprising a display, an interface unit configured to receive one or more user inputs, and processing circuitry operatively coupled to the display and the interface unit. The processing circuitry is configured to obtain an executable graph-based model comprising a first sub-graph and at least one node, wherein the first sub-graph comprises a hierarchical structure of one or more related nodes. The processing circuitry is further configured to display, on the display of the device, a graphical representation of the executable graph-based model, the graphical representation comprising a first display element associated with the first sub-graph and at least one second display element associated with the at least one node, wherein the first display element represents the first sub-graph at a first level of detail. The processing circuitry is further configured to display, on the display of the device, a first selector associated with the first display element, and receive, from the

interface unit, a first user input associated with the first selector. In response to the first user input being received, the processing circuitry is configured to replace the first display element of the graphical representation of the executable graph-based model with a third display element whilst maintaining display of the at least one second display element, wherein the third display element represents the first sub-graph at a second level of detail.

(8) According to an additional aspect of the present disclosure there is provided a method for loading of sub-graphs in executable graph-based models. The method comprises obtaining, based on a first identifier, a first manifest state, wherein the first manifest state comprises the first identifier and a second identifier, and obtaining, based on the second identifier, a second manifest state, wherein the second manifest state comprises the second identifier. The method further comprises generating a first manifest from the first manifest state, wherein the first manifest comprises a first node state storage location, the first identifier, and the second identifier. The method further comprises generating a second manifest from the second manifest state, wherein the second manifest comprises a second node state storage location and the second identifier. The method further comprises obtaining a first node state and a second node state from the first node state storage location and the second node state storage location respectively. The method further comprises generating a first node based on the first manifest and the first node state, wherein the first node comprises the first identifier and the second identifier, generating a first graph node based on the second manifest and the second node state, wherein the first graph node comprises the second identifier, and generating a first sub-graph of an executable graph-based model by associating the first node with the first graph node.

Description

BRIEF DESCRIPTION OF FIGURES

- (1) Embodiments of the present disclosure will now be described, by way of example only, and with reference to the accompanying drawings, in which:
- (2) FIG. 1 shows a prior-art n-tier architecture;
- (3) FIG. 2 shows an executable graph-based model according to an aspect of the present disclosure;
- (4) FIG. 3 shows a system for executable graph-based models according to an aspect of the present disclosure;
- (5) FIG. 4A shows the general structure of a node within an executable graph-based model according to an aspect of the present disclosure;
- (6) FIG. 4B illustrates a hyper-edge according to an embodiment of the present disclosure;
- (7) FIG. 4C shows an executable node according to an aspect of the present disclosure;
- (8) FIG. 4D shows a node having a leaf graph node type according to an aspect of the present disclosure;
- (9) FIG. 4E shows a node having a composite graph node type according to an aspect of the present disclosure;
- (10) FIG. 5A shows a graph-based model comprising multiple sub-graphs according to an aspect of the present disclosure;
- (11) FIG. 5B illustrates a traversal path for identifying a sub-graph according to embodiments of the present disclosure;
- (12) FIG. 5C illustrates a first step in the generation of an executable graph-based model from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure;
- (13) FIG. 5D illustrates a second step in the generation of an executable graph-based model from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure;
- (14) FIG. 5E illustrates a third step in the generation of an executable graph-based model from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure;

- (15) FIG. 5F illustrates a fourth step in the generation of an executable graph-based model from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure;
- (16) FIG. 5G illustrates a fifth step in the generation of an executable graph-based model from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure;
- (17) FIG. 6A illustrates the decomposition of an executable node for storage according to an embodiment of the present disclosure;
- (18) FIG. 6B shows the decomposition of a sub-graph for storage according to an aspect of the present disclosure;
- (19) FIG. 6C shows the decomposition of a shared graph relationship for persistent storage according to an embodiment of the present disclosure;
- (20) FIG. 7A shows a graph-based model with sub-graphs according to an embodiment of the present disclosure;
- (21) FIG. 7B shows an executable graph-based model generated from the graph-based model shown in FIG. 7A according to an embodiment of the present disclosure;
- (22) FIGS. 7C and 7D shows a sequence diagram of an example execution of the executable graph-based model shown in FIG. 7A;
- (23) FIG. 8A-8E illustrate the visualization of hierarchical graphs and sub-graphs according to an aspect of the present disclosure;
- (24) FIG. 9 shows a method for dynamic execution of sub-graphs within executable graph-based models according to an aspect of the present disclosure;
- (25) FIG. 10 shows a method for executing processing logic associated with overlay within a sub-graph according to an embodiment of the present disclosure;
- (26) FIG. 11 shows a method for loading and unloading one or more nodes within a sub-graph according to an embodiment of the present disclosure;
- (27) FIG. 12 shows a method for interactive visualization of sub-graphs within an executable graph-based model according to an aspect of the present disclosure;
- (28) FIG. 13 shows a method for loading of sub-graphs in executable graph-based models according to an aspect of the present disclosure;
- (29) FIG. 14 shows a method for storage management of sub-graphs in executable graph-based models according to an aspect of the present disclosure; and
- (30) FIG. 15 shows an example computing system for carrying out the methods of the present disclosure.

DETAILED DESCRIPTION

- (31) Existing architectures, such as that described in relation to FIG. 1 above, maintain a forced technical, and sometimes physical, separation between the processing logic and the data. As previously stated, the technical and physical separation of data and processing logic can be inhibitive to the types of architectural systems that can be created. Furthermore, the complexity of n-tier architectures, and their strict separation of functionality (layers), can severely impact system real-time processing performance. This, in turn, leads to processing delays or latency which reduces the applicability of such architectures being used in time-critical application settings such as medical devices, autonomous vehicles, and real-time control systems. In addition, the central storage of all data within a single database or database layer (e.g., the database layer **108** show in FIG. 1) restricts the ways in which a user may access, maintain, and manage their personal data stored by an enterprise within the single database or database layer.
- (32) The present disclosure is directed to the decomposition of executable graph-based models into any level of reusable sub-graph structure. Executable graph-based models dynamically combine data and data processing functionality at run-time whilst their separability may be maintained when at rest. This is illustrated in FIG. 2.
- (33) FIG. 2 illustrates an executable graph-based model **202** according to an aspect of the present disclosure.

(34) The executable graph-based model **202** is generally formed of a data structure (i.e., a graph-based model, or graphical model) comprising a plurality of nodes **204-08** which can be functionally extended with processing logic via the use overlays **210, 212**. Each overlay comprises processing logic, such as processing logic **214** and **216** which are associated with overlays **210** and **212** respectively. At run-time, data such as data **218, 220** is associated with nodes within the executable graph-based model **202** and the overlays **210, 212** provide the functionality to respond to stimuli an interact with, manipulate, or otherwise process the data. As such, the structure and functionality of the data processing is separate from the data itself when offline (or at rest) and is combined dynamically at run-time.

(35) As such, the executable graph-based model **202** maintains separability of the data and the data processing logic when offline thereby allowing the data owner to maintain control over their data. Moreover, by integrating the data and the data processing logic within a single model, processing delays or latency are reduced because the data and the processing logic exist within the same logical system. Therefore, the executable graph-based model **202** is applicable to a range of time-critical systems where efficient processing of stimuli is required.

(36) Whilst executable graph-based models such as that shown in FIG. **2** provide a comprehensive solution architecture to real-time in-situ processing of logic with data, they utilize a single graph to define the problem space. As will be described in more detail below, the present disclosure is directed to constructing such graphs using a composition of smaller graph structures (sub-graphs) that can then be re-used and dynamically loaded and executed.

(37) FIG. **3** shows a system **300** for execution, management, and configuration of executable graph-based models which include sub-graphs according to an aspect of the present disclosure.

(38) The system **300** comprises an executable graph-based model **302** as described in brief above in relation to FIG. **2**. The system **300** further comprises an interface module **304**, a controller module **306**, a transaction module **308**, a context module **310**, a stimuli management module **312**, a data management module **314**, an overlay management module **316**, a memory management module **318**, a storage management module **320**, a security module **322**, a visualization module **324**, an interaction module **326**, an administration module **328**, an operations module **330**, an analytics module **332**, and a graph management module **334**. FIG. **3** further shows a configuration **336**, a context **338**, data **340**, stimuli **342**, a network **344**, and an outcome **346**.

(39) The skilled person will appreciate that the present description of the system **300** is not intended to be limiting, and the system **300** can include, or interface with, further modules not expressly described herein. Moreover, the functionality of two or more of the modules can be combined within a single module. Conversely, the functionality of a single module can be split into two or more further modules which can be executed on two or more devices. The modules described below in relation to the system **300** can operate in a parallel, distributed, or networked fashion. The system **300** can be implemented in software, hardware, or a combination of both software and hardware. Examples of suitable hardware modules include, a general-purpose processor, a field programmable gate array (FPGA), and/or an application specific integrated circuit (ASIC). Software modules can be expressed in a variety of software languages such as C, C++, Java, Ruby, Visual Basic, Python, and/or other object-oriented, procedural, or other programming language.

(40) The executable graph-based model **302** corresponds to the application specific combination of data and processing functionality which is manipulated, processed, and/or otherwise handled by the other modules within the system **300**. As stated above, the structure and functionality of the data processing is separate from the data itself when offline (or at rest) and is combined dynamically at run-time. As such, different executable graph-based models are utilized for different application areas and problem domains. The skilled person will appreciate that whilst only one executable graph-based model **302** is shown in FIG. **3**, in some embodiments a system stores and maintains more than one executable graph-based model.

(41) All elements within the executable graph-based model **302** (both the data and the data processing functionality) are nodes. As will be described in more detail in relation to FIG. **4A** below, a node forms the fundamental building block of all executable graph-based models. As such, the executable graph-based model **302** comprises one or more nodes which can be dynamically generated, extended, or processed by one or more other modules within the system **300** (e.g., by the data management module **314** and/or the overlay management module **316**).

(42) The interface module **304** provides a common interface between internal components of the system **300** and/or external sources. The interface module **304** provides an application programmable interface (“API”), scripting interface, or any other suitable mechanism for interfacing externally or internally with any module of the system **300**. In the example shown in FIG. **3**, the configuration **336**, the context **338**, the data **340**, and the stimuli **342** are received by the interface module **304** of the system **300** via the network **344**. Similarly, outputs produced by the system **300**, such as the outcome **346**, are passed by the interface module **304** to the network **344** for consumption or processing by external systems. In one embodiment, the interface module **304** supports one or more messaging patterns or protocols such as the Simple Object Access protocol (SOAP), the REST protocol, and the like. The interface module **304** thus allows the system **300** to be deployed in any number of application areas, operational environments, or architecture deployments. Although not illustrated in FIG. **3**, the interface module **304** is communicatively coupled (i.e., connected either directly or indirectly) to one or more other modules or elements within the system **300** such as the controller module **306**, the context module **310**, the executable graph-based model **302** and the like. In one embodiment, the interface module **304** is communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(43) The controller module **306** handles and processes interactions and executions within the system **300**. As will be described in more detail below, stimuli (and their associated contexts) provide the basis for all interactions within the executable graph-based model **302**. Processing of such stimuli may lead to execution of processing logic associated with one or more overlays within the executable graph-based model **302**. The processing of a stimulus within the system **300** may be referred to as a system transaction. The processing and execution of stimuli (and associated overlay execution) within the system **300** is handled by the controller module **306**. The controller module **306** manages all received input stimuli (e.g., the stimuli **342**) and processes them based on a corresponding context (e.g., the context **338**). The context associated with a stimulus determines the priority that is assigned to processing the stimulus by the controller module **306**. This allows each stimulus to be configured with a level of importance and prioritization within the system **300**.

(44) The controller module **306** maintains the integrity of the modules within the system **300** before, during, and after a system transaction. The transaction module **308**, which is associated with the controller module **306**, is responsible for maintaining integrity of the system **300** through the lifecycle of a transaction. Maintaining system integrity via the controller module **306** and the transaction module **308** allows a transaction to be rolled back in the event of an expected or unexpected software or hardware fault or failure. The controller module **306** is configured to handle the processing of stimuli and transactions through architectures such as parallel processing, grid computing, priority queue techniques, and the like. In one embodiment, the controller module **306** and the transaction module **308** are communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(45) As stated briefly above, the system **300** utilizes a context-driven architecture whereby a stimulus within the system **300** is associated with a context which is used to adapt the handling or processing of the stimulus by the system **300**. The context module **310** manages the handling of contexts within the system **300** and is responsible for processing any received contexts (e.g., the context **338**) and translating the received context to an operation execution context. In some examples, the operation execution context is larger than the received context because the context

module **310** supplements the received context with further information necessary for the processing of the received context. The context module **310** passes the operational execution context to one or more other modules within the system **300** to drive the execution of the stimulus associated with the operational execution context. Contexts within the system **300** can be external or internal. While some contexts apply to all application areas and problem spaces, some applications may require specific contexts to be generated and used to process received stimuli. As will be described in more detail below, the executable graph-based model **302** is configurable (e.g., via the configuration **336**) so as only to execute within a given execution context for a given stimulus.

(46) The stimuli management module **312** processes externally received stimuli (e.g., the stimuli **342**) and any stimuli generated internally from any module within the system **300**. The stimuli management module **312** is communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302** to facilitate processing of stimuli within the executable graph-based model **302**. The system **300** utilizes different types of stimuli such as a command (e.g., a transactional request), a query, or an event received from an external system such as an Internet-of-Things (IoT) device. As previously stated, a stimulus can be either externally or internally generated. For example, a stimulus can be an event internally triggered (generated) from any of the modules within the system **300**. Such internal stimuli indicate that something has happened within the system **300** such that subsequent handling by one or more other modules within the system **300** may be required. Internal stimuli can also be triggered (generated) from execution of processing logic associated with overlays within the executable graph-based model **302**. The stimuli management module **312** communicates and receives stimuli in real-time or near-real-time. In some examples, stimuli are scheduled in a batch process. The stimuli management module **312** utilizes any suitable synchronous or asynchronous communication architectures or approaches in communicating the stimuli (along with associated information). All stimuli within the system **300** are received and processed (along with a corresponding context) by the stimuli management module **312**, which then determines the processing steps to be performed. In one embodiment, the stimuli management module **312** processes the received stimuli in accordance with a predetermined configuration (e.g., the configuration **336**) or dynamically determines what processing needs to be performed based on the contexts associated with the stimuli and/or based on the state of the executable graph-based model **302**. In some examples, processing of a stimulus results in one or more outcomes being generated (e.g., the outcome **346**). Such outcomes are either handled internally by one or more modules in the system **300** or communicated via the interface module **304** as an external outcome. In one embodiment, all stimuli and corresponding outcomes are recorded for auditing and post-processing purposes (e.g., by the operations module **330** and/or the analytics module **332**).

(47) The data management module **314** manages all data or information within the system **300** (e.g., the data **340**) for a given application. Operations performed by the data management module **314** include data loading, data unloading, data modelling, and data processing. The data management module **314** is communicatively coupled (i.e., connected either directly or indirectly) to one or more other modules within the system **300** to complete some or all of these operations. For example, data storage is handled in conjunction with the storage management module **320** (as described in more detail below).

(48) The overlay management module **316** manages all overlays within the system **300**. Operations performed by the overlay management module **316** includes overlay and overlay structure modelling, overlay logic creation and execution, and overlay loading and unloading (within the executable graph-based model **302**). The overlay management module **316** is communicatively coupled (i.e., connected either directly or indirectly) to one or more other modules within the system **300** to complete some or all of these operations. For example, overlays can be persisted in some form of physical storage using the storage management module **320** (as described in more detail below). As a further example, overlays can be compiled and preloaded into memory via the

memory management module **318** for faster run-time execution. The design and functionality of overlays is discussed in greater detail in relation to FIG. **4A** below.

(49) The memory management module **318** is configured to manage and optimize the memory usage of the system **300**. The memory management module **318** thus helps to improve the responsiveness and efficiency of the processing performed by one or more of the modules within the system **300** by optimizing the memory handling performed by these modules. The memory management module **318** uses direct memory or some form of distributed memory management architecture (e.g., a local or remote caching solution). Additionally, or alternatively, the memory management module **318** deploys multiple different types of memory management architectures and solutions. (e.g., reactive caching approaches such as lazy loading or a proactive approach such as write-through cache may be employed). These architectures and solutions are deployed in the form of a flat (single-tiered) cache or a multi-tiered caching architecture where each layer of the caching architecture can be implemented using a different caching technology or architecture solution approach. In such implementations, each cache or caching tier can be configured (e.g., by the configuration **336**) independently to the requirements for one or more of modules of the system **300**. For example, data priority and an eviction strategy, such as least-frequently-used (“LFU”) or least-recently-used (“LRU”), can be configured for all or parts of the executable graph-based model **302**. In one embodiment, the memory management module **318** is communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(50) The storage management module **320** manages the temporary or permanent storage of data within the system **300**. The storage management module **320** is any suitable low-level storage device solution (such as a file system) or any suitable high-level storage technology such as another database technology (e.g., relational database management system (RDBMS) or NoSQL database). The storage management module **320** is directly connected to the storage device upon which the relevant data is persistently stored. For example, the storage management module **320** can directly address the computer readable medium (e.g., hard disk drive, external disk drive, or the like) upon which the data is being read or written. Alternatively, the storage management module **320** is connected to the storage device via a network such as the network **344** shown in FIG. **3**. As will be described in more detail below in relation to FIGS. **13** and **14**, the storage management module **320** uses “manifests” to manage the interactions between the storage device and the modules within the system **300**. In one embodiment, the storage management module **320** is communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(51) The security module **322** manages the security of the system **300**. This includes the security at a system level and at a module level. Security is hardware related, network related, or software related, depending on the operational environment, the architecture of the deployment, or the data and information contained within the system **300**. For example, if the system is deployed with a web-accessible API (as described above in relation to the interface module **304**), then the security module **322** can enforce a hypertext transfer protocol secure (HTTPS) protocol with the necessary certification. As a further example, if the data or information received or processed by the system **300** contains Personally Identifiable Information (PII) or Protected Health Information (PHI), then the security module **322** can implement one or more layers of data protection to ensure that the PII or PHI are correctly processed and stored. In an additional example, in implementations whereby the system **300** operates on United States of America citizen medical data, the security module **322** can enforce additional protections or policies as defined by the United States Health Insurance Portability and Accountability Act (HIPAA). Similarly, if the system **300** is deployed in the European Union (EU), the security module **322** can enforce additional protections or policies to ensure that the data processed and maintained by the system **300** complies with the General Data Protection Regulation (“GDPR”). In one embodiment, the security module **322** is communicatively

coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302** thereby directly connecting security execution to the data/information in the executable graph-based model **302**. The security module **322** thus acts as a centralized coordinator working in conjunction with the data management module **314** and overlay management module **316** for managing and executing security-based overlays.

(52) The visualization module **324** and the interaction module **326** facilitate display and interaction of the executable graph-based model **302** and other parts of the system **300**. As described in more detail below in relation to FIGS. **8A-8E**, the visualization module **324** provides one or more displays, or visualizations, of the executable graph-based model **302** for review by a user of the system **300**, whilst the interaction module **326** processes user interactions (e.g., inputs, commands, etc.) with the displays, or visualizations, and/or any other module within the system **300**. The visualization module **324** and the interaction module **326** provide complex interactions capabilities such as standard two- and three-dimensional device interactions using a personal computer or mobile device and their attachable peripherals (e.g., keyboard, mouse, screen, etc.). Additionally, or alternatively, visualization module **324** and the interaction module **326** provide more advanced multi-dimensional user and visualization experiences such as virtual reality (“VR”) or augmented reality (“AR”) solutions. In one embodiment, the visualization module **324** and the interaction module **326** are communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(53) The administration module **328** manages all configurable aspects of the system **300** and the associated modules therein. Configuration is either directly embedded within the modules of the system **300** (for example, via hardware, bios, or other systems settings that are preset in the manufacturing process or software development and installation processes) or provided as dynamic configurations (e.g., via the configuration **336**). Such dynamic configurations are controllable and changeable by an end-user with the appropriate administrative privileges. In one embodiment, the degree of administrative privileges associated with an end-user are contained within a received context (e.g., the context **338**). Here, the end-user is a person connected to the administration module **328** via the interface module **304** or a system user directly connected to the administration module **328**. In one embodiment, the administration module **328** provides read-only access to all configuration settings or allows some (or all) of the configuration settings to be changed by specific user groups defined in the administration module **328** (e.g., all users associated with a user group having sufficient access privileges). In embodiments where configurations are pre-set or predetermined, the administration module **328** provides capabilities to reset or return the system **300** to its initial state or “factory settings”. In one embodiment, the administration module **328** is communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(54) The operations module **330** tracks operational metrics, module behavior, and the system **300**. Operational metrics tracked by the operations module **330** include the running status of each module, the operating performance of transactions performed, and any other associated metrics to help determine the compliance of the entire system, or any module thereof, in relation to non-functional requirements. In one embodiment, the operations module **330** is communicatively coupled (i.e., connected either directly or indirectly) to one or more overlays within the executable graph-based model **302**.

(55) The analytics module **332** performs any analytical processing required by the modules within the system **300**. The analytics module **332** processes any data embedded, or overlay contained, within the executable graph-based model **302** or created separately by the system **300** (e.g., the operation metrics produced by the operations module **330**). As such, the analytics module **332** is communicatively coupled (i.e., connected either directly or indirectly) to one or more nodes and/or one or more overlays within the executable graph-based model **302**.

(56) The graph management module **334** performs operations relating to sub-graphs within the

executable graph-based model **302**. As will be described in more detail below, the use of sub-graphs allows an executable graph-based model to be restructured into lower-level structural sub-graphs. At a general level, a sub-graph can be considered a hierarchically structured collection of nodes which are related in some manner (e.g., the nodes within the sub-graph are semantically related). Sub-graphs thus provide a mechanism for reusability of common data structures and/or functionality and can be more closely aligned to any contextualized stimuli received. The graph management module **334** manages all interactions between the executable graph-based model(s) and associated sub-graphs. The graph management module **334** controls the loading and unloading of these sub-graphs (which represent a portion of the executable graph-based model **302**) at run-time based on received stimuli. That is, sub-graphs within an executable graph-based model can be dynamically loaded, executed, and/or unloaded “on demand” thereby allowing the compute and memory requirements of the executable graph-based model to be efficiently managed (i.e., by avoiding the need for all sub-graphs to be held in memory at all times during run-time). The graph management module **334** is communicatively coupled (i.e., connected either directly or indirectly) to one or more nodes and/or one or more overlays within the executable graph-based model **302**. (57) Having now described the system **300** for executing and managing executable graph-based models, the description will now turn to the elements of an executable graph-based model; specifically, the concepts of a node and a sub-graph. Unlike conventional graph-based systems, all objects (e.g., data, overlays, etc.) within the executable graph-based model (e.g., the executable graph-based model **302**) are implemented as nodes. As will become clear, this allows executable graph-based models to be flexible, extensible, and highly configurable. Moreover, the structural organization of nodes into one or more sub-graphs provides an efficient and extensible mechanism for managing and executing executable graph-based models.

(58) FIG. 4A shows the general structure of a node **402** within an executable graph-based model, such as the executable graph-based model **302** shown in FIG. 3, according to an aspect of the present disclosure.

(59) FIG. 4A shows a node **402** which corresponds to the core structure of an executable graph-based model (e.g., the executable graph-based model **302** shown in the system **300** of FIG. 3) and which forms the foundational building block for all data and data processing logic within the executable graph-based model. The node **402** comprises properties **404**, inheritance identifiers **406**, and node type **408**. The node **402** optionally comprises one or more attributes **410**, metadata **412**, a node configuration **414**. The properties **404** of the node **402** include a unique identifier **416**, a version identifier **418**, a namespace **420**, and a name **422**. The properties **404** optionally include one or more icons **424**, one or more labels **426**, and one or more alternative identifiers **428**. The inheritance identifiers **406** of the node **402** comprise an abstract flag **430**, a leaf flag **432**, and a root flag **434**. The node configuration **414** optionally comprises one or more node configuration strategies **436** and one or more node configuration extensions **438**. FIG. 4A further shows a plurality of predetermined node types **440** which include an edge node type **442**, a role node type **444**, a leaf graph node type **446**, a composite graph node type **448**, and an overlay node type **450**. The plurality of predetermined node type **440** further include a data node type (not shown) and a value node type (not shown).

(60) The unique identifier **416** is unique for each node within an executable graph-based model. The unique identifier **416** is used to register, manage, and reference the node **402** within the system (e.g., the system **300** of FIG. 3). In some embodiments, the one or more alternative identifiers **428** are associated with the unique identifier **416** to help manage communications and connections with external systems (e.g., during configuration, sending stimuli, or receiving outcomes). The version identifier **418** of the node **402** is incremented when the node **402** undergoes transactional change. This allows the historical changes between versions of the node **402** to be tracked by modules or overlays within the system. The namespace **420** of the node **402**, along with the name **422** of the node **402**, is used to help organize nodes within the executable graph-based model. That is, the

node **402** is assigned a unique name **422** within the namespace **420** such that the name **422** of the node **402** need not be unique within the entire executable graph-based model, only within the context of the namespace **420** to which the node **402** is assigned.

(61) The node **402** optionally comprises one or more icons **424** which are used to provide a visual representation of the node **402** when visualized (e.g., by the visualization module **324** of the system **300** shown in FIG. 3). The one or more icons **424** can include icons at different resolutions and display contexts such that the visualization of the node is adapted to different display settings and contexts. The node **402** also optionally comprises one or more labels **426** which are used to override the name **422** when the node is rendered or visualized.

(62) The node **402** supports the software development feature of multiple inheritance by maintaining references (not shown) to zero or more other nodes, which then act as the base of the node **402**. This allows the behavior and functionality of a node to be extended or derived from one or more other nodes within an executable graph-based model. The inheritance identifiers **406** of the node **402** provide an indication of the inheritance-based information, which is applicable, or can be applicable, to the node **402**. The inheritance identifiers **406** comprise a set of Boolean flags which identify the inheritance structure of the node **402**. The abstract flag **430** of the inheritance identifiers **406** allows the node **402** to support the construct of abstraction. When the abstract flag **430** takes a value of “true”, the node **402** is flagged as abstract meaning that it cannot be instantiated or created within an executable graph-based model. Thus, a node having the abstract flag **430** set to “true” can only form the foundation of another node that inherits from it. By default, the abstract flag **430** of a node is set to “false”. The leaf flag **432** of the inheritance identifiers **406** is used to indicate whether any other node can inherit from the node **402**. If the leaf flag **432** is set to “true”, then no other node can inherit from the node **402** (but unlike an abstract node, a node with a leaf flag set can still be instantiated and created within an executable graph-based model). The root flag **434** of the inheritance identifiers **406** is used to indicate whether the node **402** inherits from any other node. If the root flag **434** is set to “true”, then the node **402** does not inherit from any other node. The node **402** is flagged as leaf (i.e., the leaf flag **432** is set to “true”) and/or root (i.e., the root flag **434** is set to “true”), or neither (i.e., both the leaf flag **432** and the root flag **434** are set to “false”). The skilled person will appreciate that a node cannot be flagged as both abstract and leaf (i.e., the abstract flag **430** cannot be set to “true” whilst the leaf flag **432** is set to “true”).

(63) As stated above, all elements of the executable graph-based model are defined as nodes. This functionality is in part realized due to the use of a node type. The node type **408** of the node **402** is used to extend the functionality of the node **402**. All nodes within an executable graph-based model comprise a node type which defines additional data structures and implements additional executable functionality. A node type thus comprises data structures and functionality that is common across all nodes which share that node type. The composition of a node with a node type therefore improves extensibility by allowing the generation of specialized node functionalities for specific application areas. Such extensibility is not present in prior art graph-based models. As illustrated in FIG. 4A, the node **402** and the node type **408** are one logical unit which are not separated in the context of an executing system at run-time (i.e., in the context of execution of an executable graph-based model).

(64) FIG. 4A shows the plurality of predetermined node types **440** which provides a non-exhaustive list of node types which can be associated with a node, such as the node **402**. Node types include a data node type (not shown in FIG. 4A), an edge node type **442**, a role node type **444**, a leaf graph node type **446**, a composite graph node type **448**, and an overlay node type **450**.

(65) The data node type (also referred to as a vertex or vertex node type) comprises common data structure and functionality related to the “things” modelled in the graph—i.e., the data.

(66) The edge node type **442** comprises common data structures and functionality related to joining two or more nodes. A node having the edge node type **442** can connect two or more nodes and thus the edge node type **442** constructs associations and connections between nodes (for example

objects or “things”) within the executable graph-based model. The edge node type **442** is not restricted to the number of nodes that can be associated or connected by a node having the edge node type **442**. The data structures and functionality of the edge node type **442** thus define a hyper-edge which allows two or more nodes to be connected through a defined set of roles. As will be described in more detail below, a role which defines a connective relationship involving an edge is either a (standard) role, as is known within standard hyper-graph theory such that the role merely defines a connection between the edge and another node, or the role is a node having the role node type **444**. These concepts are illustrated in FIG. **4B** described below.

(67) FIG. **4B** illustrates the concept of a hyper-edge connecting two or more nodes through a defined set of roles according to an embodiment of the present disclosure.

(68) FIG. **4B** shows a simplified representation of an edge node **452** which comprises an edge node type **454** (within the context of the example shown in FIG. **4A**, the edge node **452** corresponds to the node **402** where the node type **408** is the edge node type **442**). The edge node type **454** comprises a plurality of roles which each define a connective relationship involving the edge node **452**, e.g., a connective relationship between the edge node **452** and another node. The plurality of roles of the edge node type **454** comprises a first role node **456** and a role **458**. The plurality of roles optionally comprises a further role in the form of a second role node **460**. The first role node **456** is a node having a role node type (i.e., the role node type **444** shown in FIG. **4A**) and defines a connective relationship between the edge node **452** and a first node **462**. The role **458** defines a connective relationship between the edge node **452** and a second node **464**. The second role node **460** is a node having the role node type and defines a relationship without expressly defining the node to which the edge connects. Whilst the example in FIG. **4B** shows the edge node type **454** having two, or even three, roles, the number of roles (and thus the number of connections) that an edge node type can have is not so limited.

(69) As stated above, a role defines a connective relationship involving the edge node **452** (via the edge node type **454**) and can be either a (standard) role, such as the role **458**, or a role node, such as the first role node **456** or the second role node **460**. The standard role simply defines a connective relationship between an edge node and another node. Thus, in the example shown in FIG. **4B**, the role **458** defines the connection between the edge node **452** and the second node **464** (via the edge node type **454**). A role node is a node having a role node type (i.e., the role node type **444** shown in FIG. **4A**) and, like the (standard) role, defines a connective relationship involving an edge.

However, because a role node is a node, a role node gains the capabilities, functionality, and extensibility of a node (as described in relation to FIG. **4A**). A role node thus describes a potentially more complex connective relationship than a (standard) role. In the example shown in FIG. **4B**, the first role node **456** defines a connective relationship between the edge node **452** and the first node **462** (via the edge node type **454**). Beneficially, by utilizing the first role node **456** to define the connective relationship between the edge node **452** and the first node **462** the capabilities afforded to a node are provided to the first role node **456**. For example, and as will be described in more detail below, one or more overlay nodes can be associated with a role node to imbue the role node with processing logic thus allowing the role node to process data, respond to stimuli, etc. Moreover, a role node need not define a connective relationship to a node, as illustrated by the second role node **460**. Because the second role node **460** is itself a node, the second role node **460** encompasses the data structures and functionality of a node thereby avoiding the need to define the connecting node directly.

(70) Referring once again to FIG. **4A**, the plurality of predetermined node types **440** further comprise the leaf graph node type **446** and the composite graph node type **448**. As will be described in more detail in relation to FIG. **4D**, the leaf graph node type **446** and the composite graph node type **448** are node types which provide the mechanism and functionality for associating nodes with sub-graphs thereby allowing an executable graph-based model to be structured according to smaller units of data/functionality (i.e., sub-graphs).

(71) The overlay node type **450** is used to extend the functionality of a node, such as the node **402**, to incorporate processing logic thereby leading to the generation of an executable node, as will be described in more detail below in relation to FIG. 4C.

(72) The one or more attributes **410** correspond to the data associated with the node **402** (e.g., the data represented by the node **402** within the executable graph-based model as handled by a data management module such as the data management module **314** of the system **300** shown in FIG. 3). Because not all nodes within an executable graph-based model is associated with data, a node need not have any attributes. Each of the one or more attributes **410** are stored in any suitable format such as a data triplet of name, value type, and value.

(73) The node **402** optionally comprises metadata **412** (e.g., data stored as a name, value type, and value triplet) which is associated with either the node **402** or one or more of the one or more attributes **410** of the node **402**.

(74) The node configuration **414** provides a high degree of configurability for the different elements of a node. The node configuration **414** optionally comprises one or more node configuration strategies **436** and/or one or more node configuration extensions **438** which are complex data types. An example of a concrete node configuration strategy is an identifier strategy, associated with the configuration of the unique identifier **416** of the node **402**, which creates Snowflake identifiers. A further example of a concrete node configuration strategy is a versioning strategy, associated with the configuration of the version identifier **418** of the node **402**, which supports major and minor versioning (depending on the type of transactional change incurred by the node **402**).

(75) According to an aspect of the present disclosure, the structure and functionality of the node **402** (as described above) can be dynamically extended using the concept of an executable node. As described in relation to FIG. 4C below, an executable node provides processing functionality (i.e., processing logic) for a base node via one or more associated overlay nodes.

(76) FIG. 4C shows an executable node **466** according to an aspect of the present disclosure.

(77) The executable node **466** comprises a base node **468** and an overlay manager **470**. The overlay manager **470** registers and maintains one or more overlay nodes associated with the base node **468**, such as the first overlay node **472** and the second overlay node **474**. The first overlay node **472** has a first overlay node type **476** and the second overlay node **474** has a second overlay node type **478**.

(78) The executable node **466** is itself a node; that is, the executable node **466** extends the node **402** (or is a subtype of the node **402**) such that all the functionality and properties of the node **402** extend to the executable node **466**. The executable node **466** also dynamically extends the functionality of the base node **468** by associating the overlays maintained by the overlay manager **470** with the base node **468**. The executable node may thus be considered a composition of a base node and an overlay node and may alternatively be referred to as a node with overlay. For example, the base node **468** may have a data node type associated with a user, and the overlay manager **470** may comprise an encryption overlay which has processing logic that encrypts the attribute values of the base node **468** prior to the values being saved or output from the system. Therefore, the executable node **466** acts as a decorator of the base node **468** adding the functionality of the overlay manager **470** to the base node **468**.

(79) The skilled person will appreciate that the base node **468** refers to any suitable node within an executable graph-based model. As such, the base node **468** can be a node having a type such as a data node type, an edge node type, or the like. Alternatively, the base node **468** can itself be an executable node such that the functionality of the base (executable) node **468** is dynamically extended. In this way, complex and powerful processing functionality can be dynamically generated by associating and extending overlay nodes.

(80) The overlay manager **470** registers and maintains one or more overlay nodes associated with the base node **468**, such as the first overlay node **472** and the second overlay node **474**. The assignment of an overlay node to a base node (via the overlay manager **470**) endows the base node

with processing logic and executable functionality defined within the overlay node. Extending the functionality of a base node through one or more overlay nodes is at the heart of the dynamic generation of executable graph-based models according to an aspect of the present disclosure. As illustrated in FIG. 2 above, the data (e.g., a data node as represented by the base node **468** in FIG. 4C) and the functionality which acts upon that data (e.g., an overlay node) can be separated and independently maintained offline, but at run-time, an association between the data node and the overlay node is determined and an executable node is generated (e.g., the executable node **466** shown in FIG. 4C).

(81) An overlay node, such as the first overlay node **472** or the second overlay node **474**, is a node having an overlay node type (alternatively referred to as an overlay type) assigned to its node type. As shown in FIG. 4C, the first overlay node **472** has the first overlay node type **476** and the second overlay node **474** has the second overlay node type **478**. Different overlay node types are used to realize different functionality. Example overlay node types include an encryption overlay node type, an obfuscation overlay node type, an audit overlay node type, a prediction overlay node type, and the like. For example, if the first overlay node type **476** is an obfuscation node type and the second overlay node type **478** is an encryption node type then the functionality of the base node **468** is extended to provide obfuscation and encryption of attribute values of the base node **468**. The skilled person will appreciate that the list of overlay types is in no way exhaustive and the number of different overlay types that can be realized is not limited. Because an overlay node is itself a node, all functionality of a node described in relation to the node **402** of FIG. 4A is thus applicable to an overlay node. For example, an overlay node comprises a unique identifier, a name, etc., can have attributes (i.e., an overlay node can have its own data defined), supports multiple inheritance, and can be configured via node configurations. Furthermore, because an overlay node is a node, the overlay node can have one or more overlay nodes associated therewith (i.e., the overlay node is an overlay with overlay node). Moreover, the processing functionality of an overlay node extends to the node type of the node to which the overlay node is applied.

(82) An overlay node, such as the first overlay node **472** or the second overlay node **474**, is not bound to a single executable node or a single executable graph-based model (unlike nodes which have non-overlay node types). This allows overlay nodes to be centrally managed and reused across multiple instances of executable graph-based models.

(83) Unlike non-overlay nodes, an overlay node comprises processing logic (not shown in FIG. 4C) which determines the functionality of the overlay node. The processing logic of an overlay node comprises a block of executable code, or instructions, which carries out one or more operations. The block of executable code is pre-compiled code, code which requires interpretation at run-time, or a combination of both. Different overlay nodes provide different processing logic to realize different functionality. For example, an encryption overlay node comprises processing logic to encrypt the data (i.e., attributes) of a data node associated with the encryption overlay node, whilst an auditing overlay node comprises processing logic to record changes to the nodes state of a node associated with the auditing overlay node.

(84) The overlay manager **470** of the executable node **466** is responsible for executing all overlays registered with the overlay manager **470**. The overlay manager **470** also coordinates execution of all associated overlay nodes. In the example shown in FIG. 4C, the executable node **466** associates the base node **468** with two overlay nodes—the first overlay node **472** and the second overlay node **474**. Thus, the overlay manager **470** employs a strategy to manage the potentially cascading execution flow. Example strategies to manage the cascading execution of overlays include the visitor pattern and the pipe and filter pattern. Further examples include strategies which apply either depth-first or breadth-first processing patterns, a prioritization strategy, or a combination thereof. All execution strategies are defined and registered with the overlay manager **470** and are associated with an overlay via a node configuration extension for the overlay.

(85) As stated above in relation to FIG. 4A, a node comprises a node type which is used to extend

the functionality of the node. According to an aspect of the present disclosure, specific node types related to graphs are utilized to decompose executable graph-based models into any level of reusable sub-graph structure.

(86) FIG. 4D shows a node having a leaf graph node type according to an aspect of the present disclosure.

(87) FIG. 4D shows a leaf graph node **480** which is a node having the leaf graph node type **446**. The leaf graph node type **446** encapsulates zero or more nodes such that, in one embodiment, the leaf graph node type **446** comprises a first node **482** and a second node **484**.

(88) The leaf graph node **480** corresponds to a node, such as the node **402** shown in FIG. 4A, with a node type (i.e., node type **408**) corresponding to the leaf graph node type **446** (as shown in the plurality of predetermined node type **440** in FIG. 4A). As such, the leaf graph node **480** comprises all of the functionality and structure of the node **402** shown in FIG. 4A, but this functionality and structure is not displayed in FIG. 4D for brevity. This means that all leaf graph nodes comprise properties, identifiers, and can comprise metadata, attributes, and specific node configurations. As with any node within an executable graph-based model, leaf graph nodes can also include inheritance relationships. However, graphs also apply restrictions to the inheritance relationships that can exist within an executable graph-based model. For example, a first node within a first graph can inherit from a second node within a second graph if the first graph is a sub-graph of the second graph. However, the first node cannot inherit from the second node if the first graph is a parent graph of the second graph.

(89) The leaf graph node **480** is a leaf node, or terminal node, within the graph or sub-graph in which it exists and thus contains no further graph nodes as children. Therefore, the nodes which the leaf graph node type **446** encapsulates, such as the first node **482** or the second node **484**, are nodes which do not have a graph node type (i.e., the nodes do not have a node type of either the leaf graph node type **446** or the composite graph node type **448**). However, the first node **482** or the second node **484** may be any other type of node within the executable graph-based model such as a data node, an edge node, an executable node, etc.

(90) Because the leaf graph node **480** is a node, an executable node (or executable graph node) can be generated from the leaf graph node **480** by setting the leaf graph node **480** as the base node of an executable node (e.g., the base node **468** of the executable node **466** shown in FIG. 4C). As will be described in more detail below, when such an executable graph node is loaded, any overlays that are defined by the configuration **414** as applying to the encapsulated nodes are applied not only to the executable leaf graph node but are also cascaded down to the nodes encapsulated by the executable leaf graph node. In the example shown in FIG. 4D, if the leaf graph node **480** is employed as a base node of an executable node and encapsulates the first node **482** and the second node **484**, then when the executable node is loaded, the first node **482** and the second node **484** are generated as executable nodes (if not already so defined) and the overlays associated with the overlay manager of the executable node associated with the leaf graph node **480** are registered with the overlay managers of the executable first node and the executable second node.

(91) FIG. 4E shows a node having a composite graph node type according to an aspect of the present disclosure.

(92) FIG. 4E shows a composite graph node **486** which is a node having the composite graph node type **448**. The composite graph node type **448** encapsulates a first graph node **488**. In one embodiment, the composite graph node type **448** encapsulates a second graph node **490**. The first graph node **488** and the second graph node **490** are both nodes having either a node type of either the leaf graph node type **446** or the composite graph node type **448**.

(93) The composite graph node **486** corresponds to a node, such as the node **402** shown in FIG. 4A, with a node type (i.e., node type **408**) corresponding to the composite graph node type **448** (as shown in the plurality of predetermined node type **440** in FIG. 4A). As such, the composite graph node **486** comprises all of the functionality and structure of the node **402** shown in FIG. 4A, but

this functionality and structure is not displayed in FIG. 4E for brevity. This means that all composite graph nodes comprise properties, identifiers, and can comprise metadata, attributes, and specific node configurations. As with any node within an executable graph-based model, composite graph nodes can also include inheritance relationships. However, the restrictions to inheritance relations described above in relation to leaf graph nodes also apply to composite graph nodes.

(94) The composite graph node **486** is a composite node, or non-terminal node, within the graph or sub-graph in which it exists and thus contains one or more further graph nodes as children. Therefore, the nodes which the composite graph node type **448** encapsulates, such as the first graph node **488**, are nodes which have either the leaf graph node type **446** or the composite graph node type **448**. Although not shown in FIG. 4E, in embodiments, the composite graph node **486** also encapsulates one or more non-graph nodes as children (i.e., nodes having a node type which is not either the leaf graph node type **446** or the composite graph node type **448**).

(95) Because the composite graph node **486** is a node, an executable node (or executable graph node) can be generated from the composite graph node **486** by setting the composite graph node **486** as the base node of an executable node (e.g., the base node **468** of the executable node **466** shown in FIG. 4C). As will be described in more detail below, when such an executable graph node is loaded, overlays within the overlay manager of the executable composite graph node are applied—based on the configuration **414** as mentioned above—not only to the executable composite graph node but are also cascaded down to the graph nodes which are children of the executable composite graph node. In the example shown in FIG. 4E, if the composite graph node **486** is employed as a base node of an executable node, then when the executable node is loaded, the first graph node **488** (and optionally the second graph node **490**) is generated as an executable graph node (if not already so defined) and the overlays associated with the overlay manager of the executable node associated with the composite graph node **486** are registered with the overlay manager of the executable first graph node (and optionally the executable second graph node).

(96) Having described the structure and function of the node **402** (FIG. 4A), the executable node **466** (FIG. 4C), and the two specific graph node types (FIGS. 4D and 4E), the description will now turn to the composition of such nodes into graphs and sub-graphs.

(97) FIG. 5A shows a graph-based model **500** comprising multiple sub-graphs according to an aspect of the present disclosure.

(98) The graph-based model **500** comprises a root graph node **502**, a first leaf graph node **504**, a first composite graph node **506**, a second composite graph node **508**, a second leaf graph node **510**, and a third leaf graph node **512**. A graph-based model **500** composition cannot create a circular dependency between any of the sub-graphs. The graph-based model **500** further comprises a first node **514**, a second node **516**, a third node **518**, and a fourth node **520**. The graph-based model **500** comprises an overlay structure comprising the first overlay node **522**, a second overlay node **524**, a third overlay node **526**, a fourth overlay node **528**, and a fifth overlay node **530**.

(99) The graph-based model **500** is a configured graph or configured graph-based model. That is, the graph-based model **500** defines the data structure and overlay structure from which an executable graph-based model is generated at run-time. The graph-based model **500** comprises a graph composed of six sub-graphs, each defined according to a node within the executable graph-based model. That is, the root graph node **502**, the first leaf graph node **504**, the first composite graph node **506**, the second composite graph node **508**, the second leaf graph node **510**, and the third leaf graph node **512** all define sub-graphs.

(100) Within the context of the topology of the graph defined by the graph-based model **500**, the root graph node **502** is a composite graph node which is a root node (i.e., it has no parent node and is at the highest level within the graph). As such, the sub-graph defined by the root graph node **502** comprises the entire graph (or sub-graph) structure of the graph-based model **500** shown in FIG. 5A. The skilled person will appreciate that the graph defined by the root graph node **502** may be one of a possible plurality of sub-graphs within a larger graph-based model. In many applications, a

graph-based model, or an executable graph-based model, will comprise many sub-graphs, such as that shown in FIG. 5A. By grouping nodes into self-contained hierarchical structures (i.e., sub-graphs), these structures can be re-used, loaded on demand, and executed independently thereby improving the efficiency and extensibility of graph-based models and executable graph-based models. Moreover, each of the sub-graphs which form the graph shown in FIG. 5A (e.g., the sub-graph defined by the first leaf graph node **504**, the sub-graph defined by the second composite graph node **508**, etc.) are self-contained structures which can be loaded, unloaded, executed, etc. independently of the other sub-graphs within the overall structure.

(101) As seen in FIG. 5A, the root graph node **502** is associated with a first overlay node **522**. Because the root graph node **502** is a node within the graph-based model **500** having a composite graph node type (as described above in relation to FIG. 4E), an overlay can be associated with the node to form an executable graph node. As such, the root graph node **502** comprises a composition of a graph node (a node with a composite graph node type) and an overlay node. The root graph node **502** is a composite graph node because it comprises, via the use of a composite graph node type, the first leaf graph node **504** and the first composite graph node **506**. As such, both the first leaf graph node **504** and the first composite graph node **506** are direct children of the root graph node **502**.

(102) The first leaf graph node **504** defines a sub-graph comprising a single node—the first node **514**. The first node **514** is an executable node comprising a composition of a node and the second overlay node **524**. Within the context of the structure of the graph defined by the root graph node **502**, the first leaf graph node **504** is a leaf graph node, or terminal node, because it does not contain any child graph nodes. The first leaf graph node **504** comprises, or encapsulates, the first node **514**. As described above, the leaf graph node type provides the mechanism by which a graph node (such as the first leaf graph node **504**) can encapsulate zero or more nodes. These nodes may be owned by the sub-graph, such as is the case for the first node **514** which is owned by the first leaf graph node **504**, or the nodes may be shared across graph types as a common resource. This is exemplified by the first leaf graph node **504** which is shared by both the root graph node **502** and the second composite graph node **508**.

(103) The first composite graph node **506** is a child of the root graph node **502** and directly comprises the second composite graph node **508** and the second leaf graph node **510**. However, the sub-graph defined by the first composite graph node **506** (i.e., the sub-graph with the first composite graph node as the root node) comprises not only the directly associated nodes—the second composite graph node **508** and the second leaf graph node **510**—but also indirectly associated nodes which correspond to the child, grandchild, etc. graph nodes defined by the directly associated nodes. As such, the sub-graph having the first composite graph node **506** as its root node comprises the second composite graph node **508**, the second leaf graph node **510**, the first leaf graph node **504** (which is a child of the second composite graph node **508**), and the third leaf graph node **512** (which is also a child of the second composite graph node **508**). The first composite graph node **506** is an executable graph node because it comprises a composition of a graph node and the third overlay node **526**.

(104) The second composite graph node **508** is a child of the first composite graph node **506** and defines a sub-graph comprising the first leaf graph node **504** and the third leaf graph node **512**. The second composite graph node **508** is a composite graph node because it comprises graph nodes as children (i.e., the first leaf graph node **504**, which is shared by the root graph node **502** and the second composite graph node **508**, and the third leaf graph node **512**). The second composite graph node **508** further comprises, or encapsulates, the second node **516**.

(105) The second leaf graph node **510** is a leaf graph node within the graph because it does not contain any graph nodes as children. Rather, the second leaf graph node **510** comprises, or encapsulates, the third node **518**. The third leaf graph node **512** is also a leaf graph node within the graph because it does not contain any graph nodes as children. Rather, the third leaf graph node **512**

comprises, or encapsulates, the fourth node **520** which is an executable node associated with the fifth overlay node **530**. The third leaf graph node **512** is an executable graph node because it comprises a composition of a (leaf) graph node and the fourth overlay node **528**.

(106) The graph-based model **500** comprises an overlay structure which includes the first overlay node **522**, the second overlay node **524**, the third overlay node **526**, the fourth overlay node **528**, and the fifth overlay node **530**. Whereas the node structure described above defines the hierarchical data structure of the graphs and sub-graphs within the graph-based model **500**, the overlay structure defines the functionality, or processing logic, associated with the hierarchical data structure which is executable at run-time and able to interact with, or process, data related to the nodes within the hierarchical data structure.

(107) The hierarchical structure of the graph defined by the root graph node **502** extends to the overlay structure of the graph such that overlays within the overlay structure can be extended or overridden depending on their position within the hierarchical structure of the graph and associated configurations. As will be explained in more detail below, a child graph node can extend the overlays of its direct parent graph node; whilst the overlays of a parent graph node only extend to its direct children graph nodes. For example, because the first overlay node **522** is associated with the root graph node **502**, when the graph-based model **500** is loaded, the first overlay node **522** will be applied to the first leaf graph node **504** (which will be generated during the loading process as an executable leaf graph node within the executable graph-based model) and the first composite graph node **506** due to their having the root graph node **502** as a parent. The application of the first overlay node **522** to the first leaf graph node **504** is based on the default configuration of node behavior. In some embodiments, the node configuration (e.g., the configuration **414** of the node **402** shown in FIG. 4A) determines that an overlay such as the first overlay node **522** is not applied to the first leaf graph node **504**. This relationship between the hierarchical structure of nodes and the assignment of overlays provides an efficient mechanism for defining, maintaining, and adapting common functionality for a sub-graph. This can be particularly beneficial in situations whereby all nodes within a sub-graph should adhere to some common functionality (e.g., security functions, reporting or logging procedures, connectivity protocols, etc.) because the common functionality need only be defined once within the graph (i.e., at the root node) for it to apply, at run-time, to all nodes within the executable graph-based model.

(108) In summary, the graph-based model **500** shown in FIG. 5A corresponds to the “offline” configuration of an executable graph-based model. Representing the underlying data structure and overlay structure within a graph allows the overall model to be decomposed into reusable sub-graphs. For example, the graph-based model **500** may be an independent graph within a larger graph-based model comprising multiple other sub-graphs, nodes, etc. Moreover, sub-graphs within the graph-based model **500** can be loaded, executed, and unloaded independently of other sub-graphs within the graph-based model **500**. For example, if a stimulus associated with the first node **514** is received, then a sub-graph of the graph-based model **500** upon which the first node **514** depends (e.g., the root graph node **502** with the first overlay node **522**, the first leaf graph node **504**, and the first node **514** with the second overlay node **524**) can be loaded without requiring any non-functionally dependent graph nodes and nodes to be loaded. This leads to improvements in the use of computing resources because the portion of the sub-graph required to perform a task (e.g., process a stimulus) can be efficiently loaded and unloaded “on demand”. This is of particular benefit in time critical applications and applications involving reduced storage and reduced memory devices, where the time spent loading and unloading models, and managing the memory and storage requirements of such models, should be minimized. Therefore, system resources can be managed more efficiently by exploiting the hierarchical structure enforced by graphs to identify the portion(s) of the graph structure which are required for loading/execution based on contextualized stimuli. The identified portion(s) are then loaded and executed on demand whilst the remaining portion(s) of the graph structure remain unloaded

(109) To illustrate the above described concepts and benefits, FIGS. 5B-5F illustrate the dynamic generation of an executable graph-based model from the graph-based model **500**. The executable graph-based model is constructed from a sub-graph of the graph-based model **500** based on a stimulus (and associated context) received in relation to the fourth node **520**. As such, to process the stimulus, the sub-graph of the graph-based model **500** upon which the fourth node **520** functionally depends should be loaded into memory by generating an executable graph-based model which can then process the received stimulus.

(110) Because, in this example, the entire graph-based model **500** does not need to be loaded to process the received stimulus, the sub-graph containing the target node (the fourth node **520**) is identified and loaded. To identify the relevant sub-graph, the hierarchical structure defined by the graph-based model **500** is traversed from the target node (the fourth node **520**) to the root of the hierarchical structure (i.e., the root graph node **502**). The graph nodes within the traversal path from the root graph node **502** to the fourth node **520** correspond to the sub-graph to be loaded and from which the executable graph-based model will be generated. This traversal is only performed once when the executable graph-based model is loaded.

(111) FIG. 5B illustrates a traversal path related to the fourth node **520** of the graph-based model **500** shown in FIG. 5A according to an embodiment of the present disclosure.

(112) As stated above, the fourth node **520** corresponds to the target node—i.e., a node which is required to be loaded to perform some operation or functionality. In the present example, a stimulus associated with the fourth node **520** has been received such that the fourth node **520** needs to be loaded in order to process the stimulus. Whilst the fourth node **520** is an executable node associated with the fifth overlay node **530**, the graph-based model **500** defines a number of overlays which may be applicable to the fourth node **520**. As described above, the overlay structure of an executable graph-based model generated from a graph-based model, such as the graph-based model **500**, is dynamically generated based on the hierarchical relationship between the nodes (graph nodes, nodes, executable nodes, etc.) and the overlay nodes in the overlay structure. The traversal path is therefore used not only to identify the sub-graph containing the fourth node **520**, but also to construct the overlay structure of the executable graph-based model.

(113) The traversal from the target node (the fourth node **520**) to the root node (the root graph node **502**) is shown in FIG. 5B. The traversal path is determined by identifying the direct parent of each node and then subsequently identifying the direct parent of the parent nodes, etc. until the root graph node **502** is reached. As will be described in more detail below in relation to FIGS. 6B and 6C, parent graph nodes can be identified because child graph nodes are stored with a reference to their direct parent thereby allowing a bottom-up traversal at load time. Determining the traversal path thus comprises a number of stages, or steps, where each stage comprises a traversal from a child to a parent.

(114) The traversal starts from the fourth node **520** (i.e., the node to which the received stimulus relates) and then proceeds to the third leaf graph node **512** which encapsulates the fourth node **520**. This is shown in FIG. 5B by the jump from the circled number “1” to the circled number “2”. The traversal then proceeds to the parent of the third leaf graph node **512**, which is the second composite graph node **508**. This is shown in FIG. 5B by the jump from the circled number “2” to the circled number “3”. The traversal then proceeds to the parent of the second composite graph node **508**, which is the first composite graph node **506**. This is shown in FIG. 5B by the jump from the circled number “3” to the circled number “4”. The traversal does not proceed to the first leaf graph node **504** or the second leaf graph node **510** because these nodes are children of the second composite graph node **508** and the first composite graph node **506** respectively and thus do not fall within the direct traversal path from the fourth node **520** to the root graph node **502**. The traversal continues from the first composite graph node **506** to the root graph node **502** wherein the traversal terminates because the traversal has reached the root of the graph containing the target node (i.e., containing the fourth node **520**). This is shown in FIG. 5B by the jump from the circled number “4”

to the circled number “5”.

(115) The traversal path shown in FIG. 5B corresponds to the sub-graph of the overall graph defined by the graph-based model **500** containing the fourth node **520**. Put another way, the traversal path identifies the set of nodes and overlay nodes upon which the target node (the fourth node **520**) functionally depends. As such, the set of nodes within the traversal path corresponds to the nodes which should be loaded in order to process stimuli related to the target node. As stated above, loading the nodes within the identified sub-graph (i.e., generating an executable graph-based model from the sub-graph) requires the overlay structure of the executable graph-based model being generated to be constructed based on the overlay structure and hierarchy of the sub-graph. Beneficially, the traversal path is used in a reverse traversal order—from the root graph node **502** to the fourth node **520**—to construct the overlay structure and generate the executable-graph based model. This is illustrated in FIGS. 5C-5G.

(116) FIG. 5C illustrates a first step in the generation of an executable graph-based model **532** from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure.

(117) The root graph node **502** corresponds to the last node within the traversal path (as seen in FIG. 5B) and is thus the first node to be generated within the executable graph-based model **532** of the sub-graph. The root graph node **502** is an executable composite graph node associated with the first overlay node **522**. Consequently, an executable composite graph node is generated within the executable graph-based model **532**. The skilled person will appreciate that although the integers are the same between FIGS. 5A-5B and FIGS. 5C-5G, the root graph node **502** shown in FIG. 5C corresponds to the “run-time” root graph node **502** within the generated executable graph-based model **532**.

(118) From the traversal path, it can be inferred that the next node to construct in the executable graph-based model **532** is the first composite graph node **506** as shown in FIG. 5D.

(119) FIG. 5D illustrates a second step in the generation of the executable graph-based model **532** from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure.

(120) The first composite graph node **506** is the second graph node within the reverse traversal path shown in FIG. 5B. As such, it is the second node to be generated within the executable graph-based model **532**. The first composite graph node **506** is an executable graph node associated with the third overlay node **526** and so it is generated as an executable composite graph node within the executable graph-based model **532**.

(121) However, because the first composite graph node **506** is a direct child of the root graph node **502**, the first overlay node **522** extends to, or is overridden by, the first composite graph node **506**. This is shown in FIG. 5D by the first overlay node **522** being associated with the first composite graph node **506** due to the hierarchical relationship between the root graph node **502** (the parent) and the first composite graph node **506** (the child). In the example shown in FIG. 5D, the first overlay node **522** is applied as an additional overlay of the first composite graph node **506**. This is illustrated in FIG. 5D by the paths referenced by the circled numbers “1” and “2”. A reference to the first overlay node **522** is registered within the overlay manager of the first composite graph node **506** such that the first composite graph node **506** has two overlay nodes registered (the first overlay node **522** and the third overlay node **526**). This illustrates a difference between the graph-based model **500** and the executable graph-based model **532** generated therefrom, since the first composite graph node **506** is associated with only a single overlay node (the third overlay node **526**) in the graph-based model **500**.

(122) By default, overlays which are “cascaded down” from a parent graph node to a child graph node extend the overlays of the child graph node. That is, by default the first overlay node **522** is extended to the first composite graph node **506** by registering the first overlay node **522** as an additional overlay within the overlay manager of the first composite graph node **506**. In embodiments, this default behavior is altered via the use of node configuration thereby providing a high degree of flexibility and configuration in the application of overlays to the sub-graph overlay

structure.

(123) Whilst the first overlay node **522** is registered as an additional overlay node of the first composite graph node **506**, not all overlay nodes of a parent graph node extend the functionality of the child graph node in this way. In some examples, an overlay node of a child graph node overrides an overlay node of a parent graph node. For example, if the overlay node of the child graph node is a more restrictive version (more restrictive extension) of the overlay node of the parent graph node.

(124) FIG. 5E illustrates a third step in the generation of the executable graph-based model **532** from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure.

(125) The second composite graph node **508** is the third graph node within the reverse traversal path shown in FIG. 5B. As such, it is the third node to be generated within the executable graph-based model **532**. The second composite graph node **508** is a non-executable node within the graph-based model **500**. However, as is shown in FIG. 5D and described below, due to the cascading construction of the overlay structure of the executable graph-based model **532**, the second composite graph-node is generated as an executable composite graph-node within the executable graph-based model **532**.

(126) Although the second composite graph node **508** is not associated with any overlay nodes within the graph-based model **500**, the second composite graph node **508** “inherits” the first overlay node **522** and the third overlay node **526** because it is a direct child of the first composite graph node **506** which, within the executable graph-based model **532**, is associated with these overlay nodes. As stated above, this represents the default behavior of the application of overlays from parent graph nodes to child graph nodes which can be modified via the use of node configuration. This is illustrated by the paths referenced by the circled numbers “3” and “4” in FIG. 5E. When the sub-graph is loaded to generate the executable graph-based model **532**, the second composite graph node **508** is generated as an executable composite graph node with references to the first overlay node **522** and the third overlay node **526** being registered in its overlay manager. This is illustrated by the circled numbers “5” and “6” in FIG. 5E.

(127) The second composite graph node **508** also encapsulates the second node **516**. The second composite graph node **508** may require the second node **516** for processing (e.g., the second node **516** is a data node holding key values) and as such should be loaded along with the second composite graph node **508**. In embodiments, encapsulated graph nodes, such as the second node **516**, are loaded depending on configuration settings or contexts associated with received stimulus. The second composite graph node **508** is a non-executable node within the graph-based model **500**. However, as its encapsulating graph node (the second composite graph node **508**) is an executable graph node, the overlay nodes associated with the encapsulating graph node are applied to the second node **516**. As such, in the executable graph-based model **532**, the second node **516** is generated as an executable node with references to the first overlay node **522** and the third overlay node **526** registered within its overlay manager. This is illustrated by the circled numbers “7” and “8” in FIG. 5E.

(128) FIG. 5F illustrates a fourth step in the generation of the executable graph-based model **532** from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure.

(129) The third leaf graph node **512** is the fourth graph node within the reverse traversal path shown in FIG. 5B. As such, it is the fourth node to be generated within the executable graph-based model **532**. The third leaf graph node **512** is an executable node associated with the fourth overlay node **528** and so it is generated as an executable leaf graph node within the executable graph-based model **532**.

(130) Because the third leaf graph node **512** is a direct child of the second composite graph node **508A** (i.e., the executable composite graph node generated from the second composite graph node **508**), the overlay nodes associated with the second composite graph node **508A** are, by default, cascaded down to the third leaf graph node **512**, as illustrated by the circled numbers “9” and “10”

in FIG. 5F. Consequently, references to the first overlay node 522 and the third overlay node 526 are added to the overlay manager of the third leaf graph node 512. This is illustrated by the circled numbers “11” and “12” in FIG. 5F. Within the executable graph-based model 532, the second leaf graph node 512 is an executable leaf graph node associated with the first overlay node 522, the third overlay node 526, and the fourth overlay node 528.

(131) FIG. 5G illustrates a fifth step in the generation of the executable graph-based model 532 from the sub-graph identified in FIG. 5B according to an embodiment of the present disclosure.

(132) As the third leaf graph node 512 is the final graph node within the reverse traversal path shown in FIG. 5B, the overlay structure of the graph nodes within the executable graph-based model 532 have been correctly constructed such that the fourth node 520 (which corresponds to the target node) can now be generated. The fourth node 520 is an executable node associated with the fifth overlay node 530 and so is generated as an executable node within the executable graph-based model 532.

(133) Because the fourth node 520 is encapsulated by the third leaf graph node 512, the overlays registered with the third leaf graph node 512 are applied to the fourth node 520 (using the default configuration, as described above). Consequently, the fourth node 520 within the executable graph-based model 532 is associated with a reference to the first overlay node 522, the third overlay node 526, the fourth overlay node 528, and the fifth overlay node 530.

(134) The executable graph-based model 532 shown in FIG. 5G thus corresponds to the executable graph-based model of the sub-graph identified in FIG. 5B—i.e., the sub-graph which contains the target node (the fourth node 520)—and which is held in memory to process the stimulus associated with the target node. Beneficially, as the overlays have been correctly registered with the nodes within the executable graph-based model 532, the stimulus is processed using the overlay manager of the fourth node 520 (that is, no further traversal is necessary to determine the overlays to be executed or the ordering of execution). The overlay manager of the fourth node 520 is configured, based on the configuration of the overlays and the internal processing logic of the overlay manager, to determine the order within which the overlay nodes associated with the fourth node 520 are to be executed to process the received stimulus. One or all of the overlay nodes will be executed depending on the stimulus, associated context, and/or configuration. Moreover, as the executable graph-based model 532 is now held in memory, any further stimuli received which are associated with the executable nodes within the executable graph-based model 532 can be efficiently and effectively processed because the overlay structure of the executable graph-based model 532 is already constructed.

(135) When the executable graph-based model 532 is unloaded, the nodes are unloaded in the traversal order from the target node (the fourth node 520) to the root node (the root graph node 502). In one embodiment, all nodes within the traversal path are unloaded. In a further embodiment, one or more nodes within the traversal path are unloaded where each of the one or more nodes satisfy an unloading criterion defined by configuration. The unloading criterion includes a time-based criterion (e.g., the unloading criterion is met after a predetermined period of time has elapsed from the processing logic being executed as defined by configuration) and a criterion defined by the received stimulus (e.g., the stimulus identifies nodes to be unloaded or to remain loaded). Moreover, any nodes which are functionally dependent on nodes which are currently loaded (i.e., nodes which are held within system memory) remain loaded.

(136) In the example described above in relation to FIGS. 5B-5G, there is a single path from the root graph node 502 to the target node (the fourth node 520). In examples involving a shared sub-graph, such as that defined by the first leaf graph node 504 shown in FIG. 5A, there are multiple possible paths from the root node to the target node. For example, if the first node 514 is the target node, then the first node 514 can be reached by the path from the root graph node 502 and the first leaf graph node 504, or the path from the root graph node 502, the first composite graph node 506, the second composite graph node 508, and the first leaf graph node 504. In consequence, the first

node **514** may be associated with either the first overlay node **522** (and the second overlay node **524**) or the first overlay node **522** and the third overlay node **526** (and the second overlay node **524**) depending on which path is taken. Therefore, context is used to determine which path is taken. For example, the context associated with the received stimulus indicates that the path from the root graph node **502** to the first leaf graph node **504** is required for processing the received stimulus such that the first node is associated with the first overlay node **522** and the second overlay node **524**.

(137) Before describing an example executable graph-based model incorporating sub-graphs, the description will turn to the decomposition of an executable nodes and sub-graphs for persistent storage, as shown in FIGS. **6A** to **6C**.

(138) FIG. **6A** illustrates the decomposition of an executable node **602** for storage according to an embodiment of the present disclosure.

(139) The executable node **602** comprises a composition of a base node **604** and an overlay node **606**. The executable node **602** comprises a state **608** with an identifier **610**, the base node **604** comprises a state **612** with an identifier **614**, and the overlay node **606** comprises a state **616** with an identifier **618**. A manifest **620-624** is generated for each of the executable node **602**, the base node **604**, and the overlay node **606**. The manifest **620** associated with the executable node **602** comprises an identifier **626** and an overlay identifier **628**. The manifest **622** associated with the base node **604** comprises an identifier **630** and the manifest **624** associated with the overlay node **606** comprises an identifier **632**.

(140) The state **608** of the executable node **602** comprises all data required to reconstruct the executable node **602** (e.g., attributes, properties, etc.). The state **608** of the executable node **602** is persistently stored along with the identifier **610**. The manifest **620** is generated for the executable node **602** and comprises the identifier **626** (which is the same as the identifier **610**), the storage location of the state **608** of the executable node **602**, and the overlay identifier **628** (which is the same as the identifier **632**). The overlay identifier **628** thus identifies the manifest **624** associated with the overlay node **606**. A manifest state (not shown) is then generated for the manifest **620** and persistently stored along with the identifier **626**.

(141) The state **612** of the base node **604** comprises all data required to reconstruct the base node **604** (e.g., attributes, properties, etc.) and is persistently stored along with the identifier **614**. The manifest **622** is generated for the base node **604** and comprises the identifier **630** and the storage location of the state **612** of the base node **604**. The identifier **614** of the state **612** and the identifier **630** of the manifest **622** is the same as the identifier **610** of the state **608** of the executable node **602** (which is also the same as the identifier **626** of the manifest **620** of the executable node **602**). A manifest state (not shown) is then generated for the manifest **622** and persistently stored along with the identifier **630**. Thus, the states, manifests, and manifest states for the executable node **602** and the base node **604** all comprise the same, shared, identifier. A shared identifier can be used in this instance because the states, manifests, and manifest states are stored separately.

(142) The state **616** of the overlay node **606** comprises all data required to reconstruct the overlay node **606** (e.g., attributes, properties, processing logic, etc.) and is persistently stored along with the identifier **618**. The manifest **624** is generated for the overlay node **606** and comprises the identifier **632**, which is the same as the identifier **618**, as well as the storage location of the state **616**. A manifest state (not shown) is then generated for the manifest **624** and is persistently stored along with the identifier **632**.

(143) FIG. **6B** shows the decomposition of a sub-graph for storage according to an aspect of the present disclosure.

(144) FIG. **6B** shows a graph node **634** which comprises a state **636** and is associated with a graph identifier **638**. A manifest **640** is associated with the graph node **634** and comprises a manifest state **642**. FIG. **6B** further shows a node **644** which comprises a state **646** and a unique identifier **648**. A manifest **650** is associated with the node **644** and comprises a manifest state **652**.

(145) The state **636** of the graph node **634** comprises all data required to reconstruct the graph node **634** (e.g., attributes, properties, etc.). The state **636** is persistently stored along with the graph identifier **638** which is unique for the graph within which the graph node **634** exists. The manifest **640** is generated for the graph node **634** and comprises the graph identifier **638** as well as the storage location of the state **636** of the graph node **634**. The manifest state **642** is generated from the manifest **640** and persistently stored. The manifest state **642** also comprises the graph identifier **638** thereby allowing all the components of the graph node **634** to be reconstructed from the manifest state **642**.

(146) The node **644** is a node of the graph node **634** (i.e., either an encapsulated node or a child graph node). The state **646** of the node **644** comprises all data required to reconstruct the node **644** (e.g., attributes, properties, etc.). The state **646** is persistently stored along with the unique identifier **648** of the node **644** and the graph identifier **638**. If the node **644** is a graph node, then the unique identifier **648** is itself a graph identifier. By storing the graph identifier **638** with the node **644**, the association between the node **644** and the graph in which it resides (i.e., as defined by the graph node **634**) can be identified without requiring the graph node **634** to maintain a list of all nodes to which it is associated. The manifest **650** is generated for the node **644** and comprises the unique identifier **648** and the graph identifier **638**, as well as the storage location of the state **646** of the node **644**. The manifest state **652** is generated from the manifest **650** and persistently stored. The manifest state **652** also comprises the unique identifier **648** and the graph identifier **638** thereby allowing all the components of the node **644**, as well as its relationship to the graph in which it exists, to be reconstructed from the manifest state **652**.

(147) As shown in FIG. **6B**, each child node (e.g., the node **644**) contains a direct reference back to its parent node. As such, a graph can be loaded by finding all manifest states associated with the graph identifier and loading these nodes. As described above, the maintenance of a reference from a child (graph node) to a parent (graph node) allows executable graph-based models to be efficiently generated by allowing the sub-graph containing a node to be loaded quickly and efficiently (i.e., by following a bottom-up traversal as illustrated in FIG. **5B**).

(148) The relationship shown in FIG. **6B** holds if the node **644** is owned by the graph node **634**. If a graph is shared, then the relationship does not indicate an owner hierarchy and so the references act in the opposite way. This is illustrated in FIG. **6C**.

(149) FIG. **6C** shows the decomposition of a shared graph relationship for persistent storage according to an embodiment of the present disclosure.

(150) FIG. **6C** shows the graph node **634** shown in FIG. **6B** but in this instance, the graph node **634** uses, or shares, a second graph defined by the graph node **654**. That is, unlike in the example in FIG. **6B**, the graph node **654** is not owned by the graph node **634**. As such, the graph node **634**, the state **636**, the manifest **640**, and the manifest state **642** maintain references to the graph identifier **658** associated with the second graph as defined by the graph node **654**. Loading such a structure begins with the first graph (i.e., the graph associated with the graph identifier **638**) and then proceeds to loading all shared dependencies if they are not already loaded (i.e., the second graph associated with the graph identifier **658**). Unloading only occurs for the second graph if all other graphs which use the second graph have also been unloaded.

(151) As will be described in more detail in relation to FIGS. **13** and **14** below, an executable graph-based model which includes sub-graphs may be stored (and loaded) using the above described decomposition. Beneficially, each component is stored separately thereby allowing a user to maintain and store their data independently of the storage of the structure and functionality of the executable graph-based model.

(152) Having described the structure and function of sub-graphs and graph nodes in FIGS. **5** and **6**, an example executable graph-based model which incorporates sub-graphs will now be described to provide further understanding of the aspects described above.

(153) FIG. **7A** shows a graph-based model **700** with sub-graphs according to an embodiment of the

present disclosure.

(154) The graph-based model **700** is directed to an example quality control (QC) system within a manufacturing environment. The graph-based model **700** comprises a storage sub-graph node **702** which defines a storage sub-graph, a device sub-graph node **704** which defines a device sub-graph, and a shared sub-graph node **706** which defines a shared sub-graph. The storage sub-graph node **702** encapsulates a historical data node **708** and comprises a security overlay **710**. The device sub-graph node **704** encapsulates a protocol node **712**, a sensor node **714**, and a parameter node **716** and comprises a security overlay **718**. The sensor node **714** of the device sub-graph node **704** comprises an input overlay **720**. The shared sub-graph node **706** encapsulates a standards node **722** and a quality control (QC) edge node **724** and comprises a security overlay **726**. The QC edge node **724** comprises an alert overlay **728** and a comparison overlay **730**.

(155) The graph-based model **700** shown in FIG. 7A comprises a plurality of nodes and an overlay structure used to perform quality control of manufacture of a tympanostomy tube (alternatively referred to as a T-tube, myringotomy tube, or grommet). Tympanostomy tubes help prevent ear infection by increasing aeration in the middle ear and allow drainage of excess fluids. To help ensure successful surgery, the manufacturing process should be controlled to identify tubes which do not meet the required dimensions (e.g., due to manufacturing variations, operator error, etc.). The QC system implemented by the graph-based model **700** helps ensure that the tympanostomy tubes are manufactured within the allowable dimensions.

(156) At a general level, the QC system works by obtaining sensor readings regarding the dimensions of a manufactured tube (e.g., readings from a metrology tool such as a gauge or micrometer or readings obtained from an image of the manufactured tube). The dimensions are compared against the target dimensions as defined within a technical standard. The dimensions are also compared to historical dimensions to identify potential recurring patterns or drift which could be symptomatic of a device failure. If the dimensions of the manufactured tube do not fall within the desired range of dimensions, then an alert is issued. The alert may result in the manufactured tube being scrapped or otherwise dealt with.

(157) The skilled person will appreciate that the manufacturing system and QC system to which the graph-based model **700** relates is deliberately simplified to provide an example through which the application of executable graph-based models and sub-graphs to such systems can be described. As such, the system is not intended to provide a detailed working embodiment of a manufacturing or QC system but is rather illustrative for the purpose of explaining the functionality and benefits of sub-graphs and executable graph-based models when applied to such real-world systems.

Moreover, the graph-based model **700** is not intended to be limiting and the skilled person will appreciate that executable graph-based models can be applied to any suitable technical application area or problem domain.

(158) The storage sub-graph node **702** comprises, or encapsulates, the historical data node **708** which contains records of historical manufacturing data. For example, the recorded dimensions of all tubes manufactured over a period of time (e.g., the last day, two days, seven days, 14 days, month, etc.). The storage sub-graph node **702** also comprises a security overlay **710** (i.e., the storage sub-graph node **702** is an executable sub-graph, or a sub-graph with overlay) which defines the security policy controlling access to the data held within the storage sub-graph. The security overlay **710** restricts access to the storage sub-graph, and thus the historical data node **708**, to a predetermined set of users. For example, a first set of users may have read-write access to the data, a second set of users may have read-only access to the data, and all other users may be denied access to the data.

(159) The device sub-graph node **704** comprises, or encapsulates, nodes which are associated with, and obtain readings from, a device from which the dimensional readings are obtained. The device sub-graph node **704** comprises the protocol node **712** and the parameter node **716** which define information and data regarding communication protocols and parameters for interfacing with the

device. The sensor node **714** is an executable node having an input overlay **720** and is used to obtain readings (i.e., dimensional measurements) from the device. The device sub-graph node **704** is associated with the security overlay **718** which defines the security policy controlling access to the data and processing logic within the device sub-graph. The security overlay **718** restricts access to the device sub-graph, and thus the sensor node **714** from which sensor readings are obtained, to a predetermined set of users. For example, a first set of users may have execute access to the processing logic of the sensor node but not write access to the data generated, a second set of users may have read-only access to the data generated, all other users may be denied access to both the data and the processing logic.

(160) The shared sub-graph node **706** comprises, or encapsulates, the QC edge node **724** which connects the storage sub-graph and the device sub-graph with the standards node **722** so that the overlays associated with the QC edge node **724** can be executed. The shared sub-graph node **706** is generated to be shared by, and “bridge”, the storage sub-graph and the device sub-graph. Because the shared sub-graph node **706** is shared, it cannot be encapsulated by either the storage sub-graph node **702** or the device sub-graph node **704**. The standards node **722** comprises all the data relating to the dimensional requirements of the manufactured nodes. The QC edge node **724** comprises the comparison overlay **730** which compares the measurements obtained from the sensor node **714** with the data held in the standards node **722** and the historical data node **708**. The QC edge node **724** also comprises the alert overlay **728** which issues an alert or notification if the comparison overlay **730** determines that the acceptable dimensions have been exceeded. The shared sub-graph node **706** has a security overlay **726** which adheres to the security overlay **710** and the security overlay **718**. That is, the security policy defined by the security overlay **726** is at least as restrictive as the security policies defined by the security overlay **710** and the security overlay **718**. This is because the QC edge node **724** has potentially a different set of users who look at the combined results.

(161) FIG. 7B shows an executable graph-based model **700A** generated from the graph-based model **700** shown in FIG. 7A.

(162) The executable graph-based model **700A** corresponds to an executable version, or runtime version, of the graph-based model **700** shown in FIG. 7A. That is, the executable graph-based model **700A** comprises the data structure and overlay structure that is held in memory after the graph-based model **700** is loaded and the executable graph-based model **700A** subsequently generated.

(163) As described in detail above in relation to FIGS. 5A-5G, the overlay structure of an executable graph-based model may differ from that of the graph-based model from which the executable graph-based model is generated. In the example shown in FIG. 7B, the executable graph-based model **700A** differs from the graph-based model **700** in FIG. 7A in that the historical data node **708**, the protocol node **712**, the parameter node **716**, and the standards node **722** are generated as executable nodes within the executable graph-based model **700A**. That is, the historical data node **708A**, the protocol node **712A**, the parameter node **716A**, and the standards node **722A** in the executable graph-based model are all executable nodes having their corresponding nodes within the graph-based model **700** as the base node. As such, the historical data node **708A** is generated such that the security overlay **710** is registered in its overlay manager, the protocol node **712A** and the parameter node **716A** are generated such that the security overlay **718** is registered in their overlay managers, and the standards node **722A** is generated such that the security overlay **726** is registered in its overlay manager. The additional difference between the graph-based model **700** and the executable graph-based model **700A** is that the security overlay **718** is associated with the sensor node **714** in the executable graph-based model **700A** and the security overlay **726** is associated with the QC edge node **724** in the executable graph-based model **700A**.

(164) As stated above in relation to FIGS. 5A-5G, the construction of the overlay structure within the executable graph-based model **700A** is achieved using a traversal-based approach which is

performed once when the graph is loaded. Consequently, the respective overlay managers within the executable graph-based model **700A** hold all the requisite overlay nodes (or references thereto) such that the overlay nodes can be efficiently and quickly executed at run-time without requiring the need to re-perform the traversal, as shown in FIG. 7C.

(165) FIG. 7C shows a sequence diagram **732** of an example execution of the executable graph-based model **700A** shown in FIG. 7B.

(166) The sequence diagram **732** shown in FIG. 7C illustrates an example execution sequence involving the QC edge node **724**, the security overlay **726**, the comparison overlay **730**, the sensor node **714**, the security overlay **718**, the input overlay **720**, the historical data node **708**, the security overlay **710**, the standards node **722**, and the alert overlay **728** shown in FIG. 7B.

(167) The execution sequence starts with a stimulus and associated context being received by the QC edge node **724**. Within the context of the wider system within which the executable graph-based model functions (e.g., the system **300** shown in FIG. 3), the stimulus and corresponding context are received by the system and processed by a stimuli management module (e.g., the stimuli management module **312** of the system **300**). The stimuli management module is communicatively coupled to the relevant overlay managers of the nodes within the executable graph-based model (e.g., the overlay manager **470** shown in FIG. 4C) such that the stimuli management module is operable to pass the stimulus and context to the relevant node's overlay manager which then passes the stimulus and context to the relevant overlay(s) for processing/execution.

(168) Upon receiving the stimulus, the overlay manager of the QC edge node **724** is configured such that the security overlay **726** is identified as the first overlay associated with the QC edge node **724** to be executed (i.e., the highest priority overlay). A further stimulus **734** is fired to execute the processing logic of the security overlay **726**. The security overlay **726** processes the stimulus to determine if access to the functionality of the QC edge node **724** is to be granted based on the associated context (e.g., if the user associated with the stimulus has sufficient privileges). An outcome **736** is then provided by the security overlay **726** indicating that the stimulus received by the QC edge node **724** (and associated context) meet the security requirements of the security overlay **726**. In consequence, a stimulus **738** is fired to carry out the request associated with the stimulus. The stimulus **738** is received by the comparison overlay **730** which then executes the relevant processing logic to begin the process of performing the quality control check.

(169) The processing logic of the comparison overlay **730** fires a stimulus **740** to obtain a sensor reading from the sensor node **714**. The overlay manager of the sensor node **714** identifies the security overlay **718** as the first overlay to be executed when such a stimulus is received. As stated above, the security overlay **718** is not associated with the sensor node **714** within the graph-based model **700** but is associated with the sensor node **714** during the generation of the executable graph-based model **700A**. A stimulus **742** is fired by the sensor node **714** and received by the security overlay **718**. The security overlay **718** executes the relevant processing logic to determine if the security requirements are met (e.g., if the user associated with the original stimulus has sufficient privileges to obtain sensor readings from the device). An outcome **744** is then provided to the sensor node **714** to indicate that the security requirements are met.

(170) A stimulus **746** is then fired and received by the input overlay **720**. The stimulus **746** is associated with a request to obtain a sensor reading or measurements of the tube. That is, in response to the stimulus **746** being received, the input overlay **720** executes processing logic to obtain one or more measurements of the tube being manufactured. As stated previously, this could include obtaining measurements from a metrology device or obtaining an image of the tube and deriving measurements therefrom. Once the measurements have been obtained, an outcome **748** is provided to the sensor node **714**. The outcome **748** comprises the measurements obtained. In an embodiment, the measurements are also stored as values of one or more attributes of the sensor node **714**. An outcome **750** comprising the measurements obtained by the input overlay **720** is then

provided from the sensor node **714** to the comparison overlay **730**.

(171) The sequence diagram **732** continues in FIG. **7D**, which shows a stimulus **752** fired from the comparison overlay **730** in order to obtain historical data from the historical data node **708**. During generation of the executable graph-based model **700A**, the security overlay **710** was associated with the historical data node **708** such that the security overlay **710** is registered with the overlay manager of the historical data node **708** in the executable graph-based model **700A**. A stimulus **754** is fired and received by the security overlay **710** which determines if the security requirements are met to provide access to the historical data. An outcome **756** is provided by the security overlay **710** to the historical data node **708** to indicate that the security requirements have been met. In consequence, an outcome **758** is provided by the historical data node **708** to the comparison overlay **730**. The outcome **758** comprises the historical data requested by the comparison overlay **730**.

(172) The processing logic of the comparison overlay **730** then causes a stimulus **760** to be fired to request the standards data from the standards node **722**. As the standards node **722** and the QC edge node **724** are within the same sub-graph (i.e., the shared sub-graph node **706**), the security overlay **726** does not need to be invoked to handle the request. An outcome **762** comprising the requested historical data is provided from the standards node **722** to the comparison overlay **730**.

(173) Once all the relevant data has been collected, the comparison overlay **730** executes processing logic to determine if the measurements of the tube fall within the acceptable ranges defined by the data obtained from the standards node **722**. The processing logic also compares the measurements to the historical data obtained from the historical data node **708** to identify patterns such as potential drift which could indicate a failure in the manufacturing process. If the measurements do not fall within the acceptable ranges, or a pattern such as drift has been identified, a stimulus **764** is fired and received by the alert overlay **728**. The alert overlay **728** then executes processing logic to issue an alert indicative of the issue identified by the comparison overlay **730**.

(174) Thus, the graph-based model **700** and the executable graph-based model **700A** shown in FIGS. **7A** and **7B** provide a reusable, scalable, and efficient approach to modelling and processing data. Because the data and data processing logic is separate while offline, the two components can be managed separately by their respective owners. Because the executable graph-based model **700A** integrates the data and the data processing functionality at run-time, the model can respond efficiently to stimuli and therefore reduce the latency of the system. Moreover, the partitioning of the graph-based model **700** and the executable graph-based model **700A** into sub-graphs provides an efficient way of managing system resources. Specifically, the graph-based model **700** may form a graph within a larger executable graph-based model which models the manufacturing system such that the executable graph-based model **700A** is only generated “on-demand”. That is, the executable graph-based model **700A** is only loaded into system memory when a stimulus is received to perform a quality control check. As such, the executable graph-based model **700A** does not need to remain in system memory thus making more efficient use of system resource and reducing the security vulnerabilities associated with maintain such data and processing logic within a “live” system.

(175) The visualization of the graph-based model **700** and the executable graph-based model **700A** shown in FIGS. **7A** and **7B** correspond to visual representations of various elements within such models. As will be described in more detail below in relation to FIGS. **8A-8E**, this visual representation, and the corresponding interactive user interface functionality, provides an efficient man-machine interface for exploring the functional structure of graphs and sub-graphs within an executable graph-based model whilst maintaining the contextual information provided by the topological layout of the overall graph.

(176) FIGS. **8A-8E** illustrate the visualization of hierarchical graphs and sub-graphs according to an aspect of the present disclosure.

(177) FIG. **8A** shows a device **802** upon which a user interface **804** is displayed. A graphical

representation **805** of an executable graph-based model is displayed within the user interface **804**. The graphical representation **805** comprises a first display element **806** associated with a sub-graph, or graph, of the executable graph-based model. A first selector **808** is also displayed within the user interface **804** along with a second selector **809A** and a third selector **809B**. Although the examples shown in FIGS. **8A-8E** correspond to executable graph-based models, the overlay structure is not shown. The skilled person will appreciate that the overlay structure of the executable graph-based models can be incorporated as part of the visualizations shown in FIGS. **8A-8E** as illustrated by the example overlay structures displayed elsewhere in the present disclosure (e.g., FIG. **7B**). Moreover, the visualizations shown in FIGS. **8A-8E** are also applicable to graph-based models. In embodiments, a user, through an administration module (e.g., the administration module **328** of the system **300** shown in FIG. **3**) utilizes an application with a user interface such as the user interface **804** to generate and configure a graph-based model, add one or more overlays to the graph-based model, and then visualize the resulting executable graph-based model. As such, the skilled person will appreciate that the user interface **804** shown in FIGS. **8A-8E** correspond to a simplified representation of these steps for the purpose of illustration and explanation of the present disclosure.

(178) As shown in FIG. **8A**, the first display element **806** is represented by a circle. The skilled person will appreciate that any suitable shape can be used for the representation of display elements which include circles, squares, polygons, etc. Moreover, whilst the examples in FIG. **8A-8E** show the display elements such as the first display element **806** containing text and integer numbering, the skilled person will appreciate that this is for reference purposes only and in implementations of the present disclosure the display elements may be displayed with different text, icons, or other visual data.

(179) The first display element **806** represents the sub-graph at a first level of detail. The first level of detail corresponds to the highest level of detail of the sub-graph. That is, although the sub-graph contains further nodes and sub-graphs, these are collapsed within the first display element **806**. However, the user is provided with a visual indication that further levels of detail are viewable because the first selector **808** and the second selector **809A** indicate that the level of detail can be increased (e.g., as indicated by the “+” icon displayed on the first selector **808** and the second selector **809A**). In contrast, the third selector **809B**, which comprises a “-” icon has been disabled thus indicating that the level of detail cannot be decreased further.

(180) Both the first selector **808** and the second selector **809A** are selectable by a user (i.e., by means of a user input at or near to the selectors) and result in the level of detail of the sub-graph being increased as shown in FIG. **8B**.

(181) FIG. **8B** shows the hierarchical graph structure shown in FIG. **8A** at a second level of detail, where the second level of detail is greater than the first level of detail associated with the sub-graph in FIG. **8A**.

(182) FIG. **8B** shows the graphical representation **805** of the executable graph-based model comprising the first display element **806**, a second display element **810** associated with an edge node, a third display element **812** associated with a node and a fourth display element **814** associated with a sub-graph. The edge node, the node, and the sub-graph are all a part of the graph associated with the first display element **806**. A fourth selector **816** is also displayed proximate the fourth display element **814**.

(183) The second display element **810**, the third display element **812**, and the fourth display element **814** are displayed in a “flower petal” layout around the first display element **806**. That is, the display elements are positioned circumferentially around the first display element **806**. This provides a visual indication to the user that the edges, nodes, and sub-graphs to which these display elements relate are all a part of the graph defined by the graph node associated with the first display element **806**. This relationship is further shown by the dotted lines extending from each of the second display element **810**, the third display element **812**, and the fourth display element **814** to

the first display element **806**. The solid line extending from the second display element **810** to the third display element **812**, and the solid line extending from the second display element **810** to the fourth display element **814** indicate a connective relationship—in this instance the node is connected to the edge (via a role which is not displayed) and the sub-graph is connected to the edge (via a role which is not displayed).

(184) Selection of the first selector **808** or the third selector **809B** by a user input would result in a decrease of the level of detail and thus the user interface **804** would return to that shown in FIG. **8A**.

(185) Selection of the second selector **809A** or the fourth selector **816** would result in the level of detail increasing. That is, the sub-graph associated with the fourth display element **814** is currently displayed in a collapsed state. Because the sub-graph associated with the fourth display element **814** is a sub-graph of the graph associated with the first display element **806**, increasing the level of detail increase the level of detail of the overall graph and the sub-graph. That is, the sub-graph associated with the fourth display element **814** is shown in FIG. **8B** at a first level of detail, and selection of the second selector **809A** or the fourth selector **816** would result in this level of detail increasing, as shown in FIG. **8C**.

(186) FIG. **8C** shows the hierarchical graph structure shown in FIG. **8B** at a greater level of detail.

(187) FIG. **8C** shows the graphical representation **805** of the executable graph-based model which includes the first display element **806**, the second display element **810**, and the third display element **812**. However, as the level of detail has been increased from that of FIG. **8B**, the fourth display element **814** associated with a sub-graph has been replaced by a fifth display element which represents the sub-graph at a level of detail greater than that shown in FIG. **8B**. The compound display element (i.e., the fifth display element) comprises the fourth display element **814**, a fifth display element **818** associated with an edge of the sub-graph, a sixth display element **820** associated with an edge of the sub-graph, a seventh display element **822** associated with a node of the sub-graph, and an eighth display element **824** associated with a node of the sub-graph. The display elements are displayed in a flower petal layout around the fourth display element **814** thereby indicating their relationship to the root node represented by the fourth display element **814**. That is, the display elements are positioned circumferentially around the fourth display element **814**. As the level of detail shown in FIG. **8C** corresponds to the maximum level of detail, the second selector **809A** has been disabled. Selection of either the third selector **809B** or the fourth selector **816** by a user input would result in the user interface **804** shown in FIG. **8B** being displayed (i.e., the level of detail is decreased by one). Selection of the first selector **808** by a user input would result in the user interface **804** shown in FIG. **8A** being displayed (i.e., the entire graph structure is collapsed down to a single display element).

(188) Additionally, the connective relationship between the edge represented by the second display element **810** and the graph represented by the fourth display element **814** in FIG. **8B** is shown in greater detail in FIG. **8C**. Specifically, it can be seen in FIG. **8C** that the edge represented by the second display element **810** connects with the node associated with the eighth display element **824**. That is, the connective relationship shown in FIG. **8B** was illustrated at a level of granularity, or detail, higher than the actual relationship, since the graph associated with the fourth display element **814** is shown in a collapsed state in FIG. **8B**.

(189) FIGS. **8D** and **8E** show alternative layouts of the graph structure according to an embodiment of the present disclosure.

(190) FIG. **8D** shows a user interface **804** comprising a graphical representation **825** of the executable graph-based. The user interface **804** comprises a first display element **826** associated with a graph, a second display element **828** associated with an edge, a third display element **830** associated with a node, and a fourth display element **832** associated with a sub-graph of the graph. A first selector **834** is shown proximate the first display element **826**.

(191) The hierarchical graph structure shown in FIG. **8D** is represented at a first level of detail. As

such it corresponds to an alternate layout to that shown in FIG. 8B. By displaying the display elements associated with the elements (i.e., nodes, edges, sub-graphs, etc.) of the graph within the boundaries of the first display element **826**, the user is provided with a visual indication that the graph elements to which the display elements relate are all part of the graph associated with the first display element **826**. In an embodiment, the user interface **804** of FIG. 8D is shown after the user interface **804** of FIG. 8A after a user input related to either the first selector **808** or the second selector **809A** is received. A user input related to either the third selector **809B** or the fourth selector **836** results in an increase in level of detail as shown in FIG. 8E.

(192) FIG. 8E shows the hierarchical graph structure shown in FIG. 8D at an increased level of detail.

(193) As shown in FIG. 8E, the fourth display element **832** associated with the sub-graph has been replaced by a fifth display element **838** whilst maintaining display of the first display element **826**, the second display element **828**, and the third display element **830**. A sixth display element **840** associated with an edge of the sub-graph, a seventh display element **842** associated with an edge of the sub-graph, an eighth display element **844** associated with a node of the sub-graph, and a ninth display element **846** associated with a node of the sub-graph are all displayed within the fifth display element **838**. The fifth display element **838**, the sixth display element **840**, the seventh display element **842**, the eighth display element **844**, and the ninth display element **846** correspond to different visual representations of the fourth display element **814**, the fifth display element **818**, the sixth display element **820**, the seventh display element **822**, and the eighth display element **824** shown in FIG. 8C.

(194) By displaying these display elements within the boundaries of the display element associated with the sub-graph, the user is provided with a visual indication that the graph elements to which the display elements relate are all part of the sub-graph associated with the fifth display element **838**. Moreover, the user is provided with a clear visual indication of the boundaries of the graphs and the hierarchical relationship of the sub-graph associated with the fifth display element **838** and the graph associated with the first display element **826**. Additionally, the connective relationship between the edge represented by the second display element **828** and the graph represented by the fourth display element **832** in FIG. 8D is shown in greater detail in FIG. 8E. Specifically, it can be seen in FIG. 8E that the edge represented by the second display element **828** connects with the node associated with the ninth display element **846**. That is, the connective relationship shown in FIG. 8D was illustrated at a level of granularity, or detail, higher than the actual relationship, since the graph associated with the fourth display element **832** is shown in a collapsed state in FIG. 8D.

(195) As the level of detail shown in FIG. 8E corresponds to the maximum level of detail, the second selector **809A** has been disabled. Selection of either the third selector **809B** or the fourth selector **836** by a user input would result in the user interface **804** shown in FIG. 8D being displayed (i.e., the level of detail is decreased by one). Selection of the first selector **834** by a user input would result in the user interface **804** shown in FIG. 8A being displayed (i.e., the entire graph structure is collapsed down to a single display element).

(196) At a general level, the transitions to and from FIG. 8B to 8C and the transitions to and from FIG. 8D to 8E show that a first display element associated with a sub-graph at a first level of detail is replaced with a second display element of the sub-graph at a second level of detail whilst maintaining display of one or more further display elements associated with other elements (i.e., nodes, edges, sub-graphs, etc.) of the executable graph-based model. The change in level of detail can be an increase (e.g., the transition from FIG. 8B to 8C) or a decrease (e.g., the transition from FIG. 8C to 8B). In either scenario, the topology and layout of the visible portions of the executable graph-based model are maintained whilst the user performs an “in-place” exploration of the hierarchical graph structure. Beneficially, this helps maintain the context of the overall executable graph-based model whilst the level of detail relating to sub-graphs within the executable graph-based model are changed. Moreover, by allowing a user to expand and collapse sub-graphs in the

manner shown in FIGS. 8A-8E, more efficient use of screen real estate can be made which can be particularly beneficial on reduced resolution displays.

(197) The description will now turn to methods which operate in conjunction with the systems and functionality described above.

(198) FIG. 9 shows a method **900** for dynamic execution of sub-graphs within executable graph-based models according to an aspect of the present disclosure.

(199) The method **900** comprises the steps of obtaining **902** an executable graph-based model comprising a plurality of sub-graphs, receiving **904** a stimulus and associated context, mapping **906**, the stimulus to a first sub-graph, and causing **908** execution of processing logic based on the mapped stimulus. The method **900** also comprises the optional steps of obtaining **910** one or more execution outcomes and outputting **912** the one or more execution outcomes. In one embodiment, the method **900** is performed by a system such as the system **300** described above in relation to FIG. 3.

(200) At the step of obtaining **902**, an executable graph-based model is obtained. The executable graph-based model comprises a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs. Each sub-graph defines a hierarchical structure of related nodes. In one embodiment, the executable graph-based model is obtained from memory such that the executable graph-based model is already loaded and generated (e.g., the executable graph-based model **532** shown in FIG. 5G). Alternatively, the executable graph-based model obtained at the step of obtaining **902** is a configured executable graph-based model (an offline executable graph-based model or a graph-based model), such as the graph-based model **500** shown in FIG. 5A. That is, the executable overlay structure has not yet been loaded and generated (as described above in relation to FIGS. 5B-5G).

(201) The nodes within each sub-graph share some common relationship. For example, the nodes may be semantically related, or the nodes may be functionally related. As such, each sub-graph represents an atomic and related unit of data and/or functionality which can be re-used or loaded/executed on demand.

(202) The first sub-graph of the executable graph-based model is defined by a first graph node such that the first graph node is a root node of the first sub-graph. In the example shown in FIG. 5A, the sub-graph is defined by the root graph node **502** which is a root node of the sub-graph.

(203) In one embodiment, the first graph node is composed of one or more further nodes within the first sub-graph. The one or more further nodes comprise one or more of: a data node, a value node, and a role node.

(204) Additionally, or alternatively, the one or more further nodes comprise an edge node comprising a first connective relationship between the first graph node and the edge node. The first connective relationship is defined by a first role. The edge node may further comprise a second connective relationship between the edge node and a second node. The second connective relationship is defined by a second role. In one embodiment, the first sub-graph comprises the second node. In a further embodiment, the second node is not a part of the first sub-graph. The second node may be a graph node which defines a sub graph of the executable graph-based model.

(205) Additionally, or alternatively, the one or more further nodes comprise a second graph node. The second graph node defines a sub-graph of the first sub-graph.

(206) At the step of receiving **904**, a stimulus and a context associated with the stimulus are received.

(207) The context is associated with a target executable node of the first sub-graph. The target executable node comprises a composition of a first node of the first sub-graph and a first overlay of the overlay structure. The first overlay comprises processing logic operable to interact with the first node.

(208) At the step of mapping **906**, the stimulus is mapped to a first sub-graph of the executable graph-based model in response to the stimulus being received. The stimulus is mapped based on the

context.

(209) Consider an example of an executable graph-based model containing four sub-graphs. Each of the sub-graphs comprise a plurality of nodes including executable nodes (i.e., nodes with overlays). A stimulus and context are received and are associated with one of the four sub-graphs by virtue of the stimulus and context being associated with a target executable node within one of the four sub-graphs. Based on this association, the stimulus is mapped to the relevant sub-graph such that the processing logic within the sub-graph can be executed. Beneficially, this is done without requiring all four sub-graphs to be loaded within memory. Once the relevant sub-graph has been identified via the step of mapping **906**, that sub-graph is loaded and executed without requiring any of the other three sub-graphs to be loaded (unless there is a functional relationship between the loaded sub-graph and one of the other sub-graphs) at the step of causing **908**, execution of processing logic within the overlay structure is caused based on the mapping. The processing logic is associated with one or more nodes of the first sub-graph.

(210) The step of causing **908** execution includes at least causing execution of the processing logic of the target executable node to which the context associated with the stimulus is related. As described in more detail below in relation to FIG. **10**, the step of causing **908** execution may also include causing execution of processing logic associated with other executable nodes within the sub-graph.

(211) At the optional step of obtaining **910**, one or more outcomes of said execution of processing logic within the overlay structure are obtained. For example, data generated by the executed processing logic may be collated, or any flags or exceptions arising as a result of executing the processing logic may be obtained.

(212) At the optional step of outputting **912**, the one or more outcomes are output. The one or more outcomes may be output to an overlay within the executable graph-based model, a unit within the system which manages the executable graph-based model (e.g., the system **300** shown in FIG. **3**), or output to one or more systems or devices external to the system.

(213) FIG. **10** shows a method **1000** for executing processing logic associated with overlay within a sub-graph according to an embodiment of the present disclosure.

(214) The method **1000** comprises the steps of performing **1002** a traversal to determine a traversal path, identifying **1004** a plurality of executable nodes, and causing **1006** execution of processing logic. In one embodiment, the method **1000** is performed after the step of receiving **904** in the method **900** and prior to the step in the method **900** of causing **908** execution such that the step of causing **1006** execution in the method **1000** corresponds to the step of causing **908** execution in the method **900**. That is, the method **1000** is performed as part of loading and generating an executable graph-based model in response to a stimulus being received.

(215) At the step of performing **1002**, a traversal is performed from the target executable node to the root node of the first sub-graph thereby determining a traversal path. This is illustrated by the traversal path shown in FIG. **5B**. A stimulus is mapped to a target executable node—the fourth node **520**—and a traversal path is determined from the target executable node to the root node—the root graph node **502**. As a result of the traversing performed at this step, the traversal path comprises a first set of nodes from the root node to the target executable node. In the example shown in FIG. **5B**, the traversal path comprises the fourth node **520**, the third leaf graph node **512**, the second composite graph node **508**, the first composite graph node **506**, and the root graph node **502**. As such, the first set of nodes includes the target executable node to which the received stimulus relates.

(216) At the step of identifying **1004**, a plurality of executable nodes are identified within the first set of nodes. Each of the plurality of executable nodes comprise a composition of a respective node of the first sub-graph and at least one overlay of the overlay structure.

(217) At the step of causing **1006**, execution of processing logic associated with each overlay of the plurality of executable nodes is caused. In the example shown in FIG. **5G**, causing execution of

processing logic associated with the fourth node **520** (the target node) would result in execution of processing logic for the first overlay node **522**, the third overlay node **526**, the fourth overlay node **528** and/or the fifth overlay node **530**.

(218) FIG. **11** shows a method **1100** for loading and unloading one or more nodes within a sub-graph according to an embodiment of the present disclosure.

(219) The method **1100** comprises the steps of performing **1102** a traversal to determine a traversal path and loading **1104** a first set of nodes. The method **1100** further comprises the optional steps of generating **1106** executable nodes as part of the step of loading **1104** the first set of nodes and unloading **1108** one or more nodes. In one embodiment, the method **1000** is performed after the step of receiving **904** in the method **900** and prior to the step in the method **900** of causing **908** execution.

(220) The step of performing **1102** a traversal corresponds to the step of performing **1002** a traversal as described above in relation to FIG. **10**.

(221) At the step of loading **1104**, the first set of nodes identified as a result of performing **1102** the traversal are loaded into a memory (e.g., a memory operatively coupled to the processing circuitry which is carrying out the method **1100**). The first set of nodes are loaded in order from the root node to the target executable node. In the example traversal path shown in FIG. **5B**, the nodes would be loaded in the order of the root graph node **502**, the first composite graph node **506**, the second composite graph node **508**, the third leaf graph node **512**, and the fourth node **520**.

(222) In an embodiment, the step of loading **1104** comprises the step of generating **1106** an executable node. When the first set of nodes comprises a node associated with at least one overlay within the overlay structure the step of generating **1106** comprises generating an executable node comprising a composition of the node and the at least one overlay. When the graph node is associated with at least one overlay (i.e., the graph node, or root graph node, is an executable node), a first set of executable nodes are generated at the step of generating **1106**. Each executable node within the first set of executable nodes comprises a composition of a respective node of the first set of nodes and the at least one overlay. For example, and as described in more detail above in relation to FIGS. **5C-5G**, when the fourth node **520** is generated, it is associated with the overlay nodes within the sub-graph such that the fourth node **520** is associated with the first overlay node **522**, the third overlay node **526**, the fourth overlay node **528**, and the fifth overlay node **530**.

(223) Although the method **1100** in FIG. **11** is shown as proceeding directly from the step of loading **1104** to the step of unloading **1108**, the skilled person will appreciate that a number of steps may occur between these steps. For example, the loaded nodes may be manipulated, executed, or otherwise used. As such, the step of unloading **1108** occurs after the step of causing **908** execution in the method **900**.

(224) At the optional step of unloading **1108**, one or more nodes within the first set of nodes are unloaded (e.g., from the memory). The one or more nodes are unloaded in order from the target executable node to the root node (i.e., in the traversal order determined at the step of performing **1102**). In one embodiment, all the nodes within the first set of nodes (i.e., the first sub-graph) are unloaded at the step of unloading **1108**. Alternatively, the one or more nodes are unloaded when an unloading criterion is met. That is, a node within the first set of nodes is unloaded when the node meets the unloading criterion. In an embodiment, the unloading criterion is met after a predetermined period of time has elapsed from the step of causing **908** execution in the method **900**. In an alternative embodiment, the unloading criterion is based on the received stimulus (i.e., the stimulus received at the receiving **904** step of the method **900** determines when and how the one or more nodes are unloaded).

(225) FIG. **12** shows a method **1200** for interactive visualization of sub-graphs within an executable graph-based model according to an aspect of the present disclosure.

(226) The method **1200** comprises the steps of obtaining **1202** an executable graph-based model comprising a first sub-graph, displaying **1204** a graphical representation of the executable graph-

based model including a first display element and a second display element, displaying **1206** a first selector, receiving **1208** a user input, and replacing **1210** the first display element with a third display element. In one embodiment, the method **1200** is performed by a device comprising a display and a user input mechanism such as that described in relation to FIG. **15** below. In an alternative embodiment, the method **1200** is performed by a system such as the system **300** described above in relation to FIG. **3**.

(227) At the step of obtaining **1202**, an executable graph-based model is obtained. The executable graph-based model comprises a first sub-graph and at least one node. The first sub-graph comprises a hierarchical structure of one or more related nodes.

(228) At the step of displaying **1204**, a graphical representation of the executable graph-based model is displayed (e.g., the graphical representation **805** of the executable graph-based model shown in FIG. **8B**). The graphical representation comprises a first display element associated with the first sub-graph (e.g., the fourth display element **814** shown in FIG. **8B**) and at least one second display element associated with the at least one node (e.g., the third display element **812** shown in FIG. **8B**). The first display element represents the first sub-graph at a first level of detail.

(229) At the step of displaying **1206**, a first selector associated with the first display element is displayed (e.g., the fourth selector **816** associated with the fourth display element **814** in FIG. **8B**). The first selector is displayed proximate the first display element. Alternatively, the first selector and the first display element are contiguous, as shown in FIG. **8B** where the fourth selector **816** and the fourth display element **814** are contiguous (i.e., their outer edges are touching).

(230) At the step of receiving **1208**, a first user input associated with the first selector is received.

(231) At the step of replacing **1210**, the first display element of the graphical representation of the executable graph-based model is replaced with a third display element in response to the first user input being received (e.g., the fourth display element **814** shown in FIG. **8B** is replaced in FIG. **8C** by the compound display element comprising the fourth display element **814**, the fifth display element **818**, the sixth display element **820**, the seventh display element **822**, and the eighth display element **824**). The third display element represents the first sub-graph at a second level of detail. As shown in the transition from the user interface **804** in FIG. **8B** to the user interface **804** in FIG. **8C** the second level of detail is greater than the first level of detail. Alternatively, the second level of detail is less than the first level of detail. For example, the transition from the user interface **804** in FIG. **8C** to the user interface **804** in FIG. **8B** where the compound display element is replaced by the fourth display element **814**. The first display element is replaced with the third display element whilst maintaining display of the at least one second display element (e.g., the third display element **812** is maintained in both FIG. **8B** at the first level of detail and in FIG. **8C** at the second level of detail).

(232) FIG. **13** shows a method **1300** for loading of sub-graphs in executable graph-based models according to an aspect of the present disclosure.

(233) The method **1300** comprises the steps of obtaining **1302** a first manifest state, obtaining **1304** a second manifest state, generating **1306** a first manifest, generating **1308** a second manifest, obtaining **1310** node states, generating **1312** a first node, generating **1314** a first graph node, and generating **1316** a first sub-graph. In one embodiment, the method **900** is performed by a system such as the system **300** described above in relation to FIG. **3**.

(234) At the step of obtaining **1302**, a first manifest state is obtained based on a first identifier (e.g., the manifest state **652** associated with the unique identifier **648** shown in FIG. **6B**). The first manifest state comprises the first identifier and a second identifier (e.g., the graph identifier **638** shown in FIG. **6B**).

(235) The first identifier is the unique identifier of the node to which the first manifest state relates. As such, the method **1300** may further comprise the step of obtaining the first identifier (not shown). For example, as a result of a user request or as a result of a stimulus or event such as a stimulus associated with dynamically loading a sub-graph. The second identifier is used to identify

the immediate, or direct, parent of the node to which the first manifest relates (as shown in FIG. 6B).

(236) At the step of obtaining **1304**, a second manifest state is obtained based on the second identifier (e.g., the manifest state **642** shown in FIG. 6B is obtained based on the graph identifier **638** obtained from the manifest state **652**). The second manifest state comprises the second identifier.

(237) At the step of generating **1306**, a first manifest is generated from the first manifest state (e.g., the manifest **650** is generated, or extracted, from the manifest state **652**). The first manifest comprises a first node state storage location, the first identifier, and the second identifier. Because the node state is stored separately from the manifest state, the location of the node state is maintained within the manifest. Here, the storage location corresponds to a file path, or location (e.g., URL), at which the node state has been persistently stored (saved).

(238) At the step of generating **1308**, a second manifest is generated from the second manifest state (e.g., the manifest **640** is generated, or extracted, from the manifest state **642**). The second manifest comprises a second node state storage location and the second identifier.

(239) At the step of obtaining **1310**, a first node state and a second node state are obtained from the first node state storage location and the second node state storage location respectively. A node state, or simply state, comprises all data required to reconstruct the node (e.g., attributes, properties, etc.). In the example shown in FIG. 6B, the state **656** is obtained from the storage location identified within the manifest **650** whilst the state **636** is obtained from the storage location identified within the manifest **640**.

(240) At the step of generating **1312**, a first node is generated based on the first manifest and the first node state (e.g., the node **644** is generated based on the manifest **650** and the state **646**). The first node comprises the first identifier and the second identifier.

(241) At the step of generating **1314**, a first graph node is generated based on the second manifest and the second node state (e.g., the graph node **634** is generated based on the manifest **640** and the state **636**). The first graph node comprises the second identifier.

(242) At the step of generating **1316**, a first sub-graph of an executable graph-based model is generated by associating the first node with the first graph node.

(243) The above steps are repeated for all parent nodes within the sub-graph until the root node of the sub-graph is reached. That is, by starting from the target node to be loaded (i.e., the first node) and recursively loading any directly referenced parent nodes, a traversal based loading is performed from the target node to the root node. For example, if the first graph node contains references to a parent node (in the way in which the first node references the first graph node via reference to the parent node's graph identifier), then the parent node is loaded from its manifest state as described above. Moreover, whilst the above is related to loading a single node and associating that node with its containing graph node, the skilled person will appreciate that the steps related to loading and generating the first node above can be repeated for one or more further nodes which are identified as part of the sub-graph by the relevant graph identifier which is stored with the node manifests. As stated in more detail in relation to FIGS. 5C-5G, the overlay structure of the executable graph-based model is then generated by "cascading down" overlay nodes from parent graph nodes to child graph nodes.

(244) In one embodiment, a shared sub-graph is loaded. The shared sub-graph is identified by a third identifier which is stored with the second manifest state. In such an embodiment, the shared sub-graph is loaded by loading the shared sub-graph based on the third identifier (i.e., using the loading method described above) and then associating the shared sub-graph with the first sub-graph. As described above, the overlays associated with executable nodes within shared sub-graphs are based on the calling path in the hierarchy which can depend on the context of the associated request or stimulus.

(245) The skilled person will appreciate that referenced overlay(s) of an executable node are not

considered in the persistence, only those that are owned by the node directly. This is because the configuration of a parent node can change, and so hard coded associations that are liable to break are avoided.

(246) FIG. **14** shows a method **1400** for storage management of sub-graphs in executable graph-based models according to an aspect of the present disclosure.

(247) The method **1400** comprises the steps of obtaining **1402** a first sub-graph, extracting **1404** a first node state, extracting **1406** a second node state, determining **1408** storage locations, generating **1410** a first sub-graph manifest, generating **1412** a first node manifest, generating **1414** a first manifest state, generating **1416** a second manifest state, and storing **1418** the manifest states. In one embodiment, the method **900** is performed by a system such as the system **300** described above in relation to FIG. **3**.

(248) At the step of obtaining **1402**, a first sub-graph of an executable graph-based model is obtained. The first sub-graph comprises a first graph node and a second node (e.g., the graph node **634** and the node **644** shown in FIG. **6B**). The first graph node defines the first sub-graph and is associated with a first identifier (e.g., the graph identifier **638** shown in FIG. **6B**).

(249) At the step of extracting **1404**, a first node state is extracted from first graph node (e.g., the state **636** shown in FIG. **6B**). The first node state comprises the first identifier. A node state, or simply state, comprises all data required to reconstruct the graph node (e.g., attributes, properties, etc.).

(250) At the step of extracting **1406**, a second node state is extracted from the second node (e.g., the state **646** shown in FIG. **6B**). The second node state comprises the first identifier and a second identifier associated with the second node (e.g., the graph identifier **638** and the unique identifier **648** shown in FIG. **6B**).

(251) At the step of determining **1408**, a storage location for the node state of each of the first graph node and the second node is determined. Because the node states are stored separately from the node manifests, the storage locations for the node states need to be determined. Here, determining a storage location corresponds to identifying or determining the file path, or location (e.g., URL), at which the node state will be, or has been, saved.

(252) At the step of generating **1410**, a first sub-graph manifest associated with the first graph node is generated (e.g., the manifest **640** shown in FIG. **6B**). The first sub-graph manifest comprises the first identifier and the storage location for the first node state. Storing the storage location of the first node state with the manifest allows the node state to be efficiently obtained when loading the sub-graph (as described above in relation to FIG. **13**).

(253) At the step of generating **1412**, a first node manifest associated with the second node is generated (e.g., the manifest **650** shown in FIG. **6B**). The first node manifest comprises the first identifier, the second identifier, and the storage location for the second node state.

(254) At the step of generating **1414**, a first manifest state for the first sub-graph manifest is generated (e.g., the manifest state **642** shown in FIG. **6B**). The first manifest state comprises the first identifier.

(255) At the step of generating **1416**, a second manifest state for the first node manifest is generated (e.g., the manifest state **652** shown in FIG. **6B**). The second manifest state comprises the first identifier and the second identifier (e.g., the graph identifier **638** and the unique identifier **648**). By storing the manifest state of the first node with both the first identifier (graph identifier) and the second identifier (unique node identifier), the sub-graph within which the first node is contained can be loaded in a bottom-up approach as described above in relation to FIG. **14**.

(256) At the step of storing **1418**, the first manifest state and the second manifest state are stored. Optionally, the first node state and the second node state are also stored at the step of storing **1418**. Beneficially, the manifest states (i.e., the data structure) and the node states (i.e., the data stored within the data structure) are stored separately thereby allowing a user to maintain and store their data independently of the storage of the structure and functionality of the executable graph-based

model. In one embodiment, the storage location of the manifest states is at a first device and the storage location of the node states are at a second, different, device. For example, the first device could be an enterprise facility managed by a company offering a service involving the executable graph-based model within which the executable node is contained, and the second device could be a device of a user of the executable graph-based model. As such, the user is able to maintain their personal data separately from, and independently to, the enterprise facility managed by the company.

(257) FIG. 15 shows an example computing system for carrying out the methods of the present disclosure. Specifically, FIG. 15 shows a block diagram of an embodiment of a computing system according to example embodiments of the present disclosure.

(258) Computing system **1500** can be configured to perform any of the operations disclosed herein such as, for example, any of the operations discussed with reference to the functional modules described in relation to FIG. 1A. The computing system **1500** can be implemented as a conventional computer system, an embedded controller, a laptop, a server, a mobile device, a smartphone, a set-top box, a kiosk, a vehicular information system, one or more processors associated with a television, a customized machine, any other hardware platform, or any combination or multiplicity thereof. In one embodiment, the computing system **1500** is a distributed system configured to function using multiple computing machines interconnected via a data network or bus system.

(259) The computing system **1500** includes one or more computing device(s) **1502**. The one or more computing device(s) **1502** of computing system **1500** comprise one or more processors **1504** and memory **1506**. One or more processors **1504** can be any general purpose processor(s) configured to execute a set of instructions. For example, one or more processors **1504** can be a processor core, a multiprocessor, a reconfigurable processor, a microcontroller, a digital signal processor (“DSP”), an application-specific integrated circuit (“ASIC”), a graphics processing unit (“GPU”), a neural processing unit (“NPU”), an accelerated processing unit (“APU”), a brain processing unit (“BPU”), a data processing unit (“DPU”), a holographic processing unit (“HPU”), an intelligent processing unit (“IPU”), a microprocessor/microcontroller unit (“MPU/MCU”), a radio processing unit (“RPU”), a tensor processing unit (“TPU”), a vector processing unit (“VPU”), a wearable processing unit (“WPU”), a field programmable gate array (“FPGA”), a programmable logic device (“PLD”), a controller, a state machine, gated logic, discrete hardware component, any other processing unit, or any combination or multiplicity thereof. In one embodiment, one or more processors **1504** include one processor. Alternatively, one or more processors **1504** include a plurality of processors that are operatively connected. For example, the one or more processors **1504** can be multiple processing units, a single processing core, multiple processing cores, special purpose processing cores, co-processors, or any combination thereof. One or more processors **1504** are communicatively coupled to memory **1506** via address bus **1508**, control bus **1510**, and data bus **1512**.

(260) Memory **1506** can include non-volatile memories such as read-only memory (“ROM”), programmable read-only memory (“PROM”), erasable programmable read-only memory (“EPROM”), flash memory, or any other device capable of storing program instructions or data with or without applied power. The memory **1506** can also include volatile memories, such as random-access memory (“RAM”), static random-access memory (“SRAM”), dynamic random-access memory (“DRAM”), and synchronous dynamic random-access memory (“SDRAM”). The memory **1506** can comprise single or multiple memory modules. While the memory **1506** is depicted as part of the one or more computing device(s) **1502**, the skilled person will recognize that the memory **1506** can be separate from the one or more computing device(s) **1502**.

(261) Memory **1506** can store information that can be accessed by one or more processors **1504**. For instance, memory **1506** (e.g., one or more non-transitory computer-readable storage mediums, memory devices) can include computer-readable instructions (not shown) that can be executed by

one or more processors **1504**. The computer-readable instructions can be software written in any suitable programming language or can be implemented in hardware. Additionally, or alternatively, the computer-readable instructions can be executed in logically and/or virtually separate threads on one or more processors **1504**. For example, memory **1506** can store instructions (not shown) that when executed by one or more processors **1504** cause one or more processors **1504** to perform operations such as any of the operations and functions for which computing system **1500** is configured, as described herein. In addition, or alternatively, memory **1506** can store data (not shown) that can be obtained, received, accessed, written, manipulated, created, and/or stored. The data can include, for instance, the data and/or information described herein in relation to FIGS. **1** to **14**. In some implementations, the one or more computing device(s) **1502** can obtain from and/or store data in one or more memory device(s) that are remote from the computing system **1500**.

(262) The one or more computing device(s) **1502** further comprise I/O interface **1514** communicatively coupled to address bus **1508**, control bus **1510**, and data bus **1512**. The I/O interface **1514** is configured to couple to one or more external devices (e.g., to receive and send data from/to one or more external devices). Such external devices, along with the various internal devices, may also be known as peripheral devices. The I/O interface **1514** may include both electrical and physical connections for operably coupling the various peripheral devices to the one or more computing device(s) **1502**. The I/O interface **1514** may be configured to communicate data, addresses, and control signals between the peripheral devices and the one or more computing device(s) **1502**. The I/O interface **1514** may be configured to implement any standard interface, such as a small computer system interface (“SCSI”), serial-attached SCSI (“SAS”), fiber channel, peripheral component interconnect (“PCI”), PCI express (“PCIe”), serial bus, parallel bus, advanced technology attachment (“ATA”), serialATA (“SATA”), universal serial bus (“USB”), Thunderbolt, FireWire, various video buses, and the like. The I/O interface **1514** is configured to implement only one interface or bus technology. Alternatively, the I/O interface **1514** is configured to implement multiple interfaces or bus technologies. The I/O interface **1514** may include one or more buffers for buffering transmissions between one or more external devices, internal devices, the one or more computing device(s), or the one or more processors **1504**. The I/O interface **1514** may couple the one or more computing device(s) **1502** to various input devices, including mice, touch screens, scanners, biometric readers, electronic digitizers, sensors, receivers, touchpads, trackballs, cameras, microphones, keyboards, any other pointing devices, or any combinations thereof. The I/O interface **1514** may couple the one or more computing device(s) **1502** to various output devices, including video displays, speakers, printers, projectors, tactile feedback devices, automation control, robotic components, actuators, motors, fans, solenoids, valves, pumps, transmitters, signal emitters, lights, and so forth.

(263) Computing system **1500** further comprises storage unit **1516**, network interface **1518**, input controller **1520**, and output controller **1522**. Storage unit **1516**, network interface **1518**, input controller **1520**, and output controller **1522** are communicatively coupled to the central control unit (i.e., the memory **1506**, the address bus **1508**, the control bus **1510**, and the data bus **1512**) via I/O interface **1514**. The network interface **1518** communicatively couples the computing system **1500** to one or more networks such as wide area networks (“WAN”), local area networks (“LAN”), intranets, the Internet, wireless access networks, wired networks, mobile networks, telephone networks, optical networks, or combinations thereof. The network interface **1518** may facilitate communication with packet switched networks or circuit switched networks which use any topology and may use any communication protocol. Communication links within the network may involve various digital or analog communication media such as fiber optic cables, free-space optics, waveguides, electrical conductors, wireless links, antennas, radio-frequency communications, and so forth.

(264) Storage unit **1516** is a computer readable medium, preferably a non-transitory computer readable medium, comprising one or more programs, the one or more programs comprising

instructions which when executed by the one or more processors **1504** cause computing system **1500** to perform the method steps of the present disclosure. Alternatively, storage unit **1516** is a transitory computer readable medium. Storage unit **1516** can include a hard disk, a floppy disk, a compact disc read-only memory (“CD-ROM”), a digital versatile disc (“DVD”), a Blu-ray disc, a magnetic tape, a flash memory, another non-volatile memory device, a solid-state drive (“SSD”), any magnetic storage device, any optical storage device, any electrical storage device, any semiconductor storage device, any physical-based storage device, any other data storage device, or any combination or multiplicity thereof. In one embodiment, the storage unit **1516** stores one or more operating systems, application programs, program modules, data, or any other information. The storage unit **1516** is part of the one or more computing device(s) **1502**. Alternatively, the storage unit **1516** is part of one or more other computing machines that are in communication with the one or more computing device(s) **1502**, such as servers, database servers, cloud storage, network attached storage, and so forth.

(265) Moreover, for example, the present technology/system may achieve the following configurations:

(266) 1. A method for dynamic execution of sub-graphs within executable graph-based models, the method comprising: obtaining, by processing circuitry, an executable graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein each sub-graph defines a hierarchical structure of related nodes; receiving, by the processing circuitry, a stimulus and a context associated with the stimulus; in response to the stimulus being received, mapping, by the processing circuitry and based on the context, the stimulus to a first sub-graph of the executable graph-based model; and causing, by the processing circuitry, execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with one or more nodes of the first sub-graph.

(267) 2. The method of 1 wherein the context is associated with a target executable node of the first sub-graph, the target executable node comprising a composition of a first node of the first sub-graph and a first overlay of the overlay structure, the first overlay comprising processing logic operable to interact with the first node such that the step of causing execution of processing logic within the overlay structure comprises causing execution of said processing logic of the first overlay.

(268) 3. The method of 2 wherein the first sub-graph of the executable graph-based model is defined by a first graph node such that the first graph node is a root node of the first sub-graph.

(269) 4. The method of 3 further comprising, in response to the stimulus being received and prior to the step of causing execution: traversing, by the processing circuitry, from the target executable node to the root node of the first sub-graph thereby determining a traversal path comprising a first set of nodes from the root node to the target executable node, wherein the first set of nodes includes the target executable node.

(270) 5. The method of 4 wherein the step of causing execution comprises: identifying, by the processing circuitry, a plurality of executable nodes within the first set of nodes, each of the plurality of executable nodes comprising a composition of a respective node of the first sub-graph and at least one overlay of the overlay structure; and causing, by the processing circuitry, execution of processing logic associated with each overlay of the plurality of executable nodes.

(271) 6. The method of 4 further comprising, after the step of traversing and prior to the step of causing execution: loading, by the processing circuitry and to memory operatively coupled to the processing circuitry, the first set of nodes of the traversal path, wherein the first set of nodes are loaded in order from the root node to the target executable node.

(272) 7. The method of 6 wherein the first set of nodes comprises a node associated with at least one overlay within the overlay structure such that the step of loading the first set of nodes comprises: generating, by the processing circuitry, an executable node comprising a composition of

the node and the at least one overlay.

(273) 8. The method of 6 wherein the first graph node is associated with at least one overlay within the overlay structure such that the first graph node is an executable node.

(274) 9. The method of 8 wherein the step of loading the first set of nodes comprises: generating, by the processing circuitry, a first set of executable nodes based on the first set of nodes and the at least one overlay, wherein each executable node within the first set of executable nodes comprises a composition of a respective node of the first set of nodes and the at least one overlay.

(275) 10. The method of 6 further comprising, after the step of causing execution: unloading, by the processing circuitry and from memory operatively coupled to the processing circuitry, one or more nodes within the first set of nodes, wherein the one or more nodes are unloaded in order from the target executable node to the root node.

(276) 11. The method of 10 wherein the one or more nodes are unloaded when an unloading criterion is met.

(277) 12. The method of 11 wherein the unloading criterion is met after a predetermined period of time has elapsed from the step of causing execution.

(278) 13. The method of 11 wherein the unloading criterion is based on the stimulus.

(279) 14. The method of 2 wherein the first graph node is composed of one or more further nodes within the first sub-graph.

(280) 15. The method of 14 wherein the one or more further nodes comprise one or more of: a data node, a value node, and a role node.

(281) 16. The method of 14 wherein the one or more further nodes comprise an edge node comprising a first connective relationship between the first graph node and the edge node.

(282) 17. The method of 16 wherein the first connective relationship is defined by a first role.

(283) 18. The method of 16 wherein the edge node further comprises a second connective relationship between the edge node and a second node.

(284) 19. The method of 18 wherein the first sub-graph comprises the second node.

(285) 20. The method of 18 wherein the second node is not a part of the first sub-graph.

(286) 21. The method of 18 wherein the second node is a graph node which defines a sub-graph of the executable graph-based model.

(287) 22. The method of 18 wherein the second connective relationship is defined by a second role.

(288) 23. The method of 14 wherein the one or more further nodes comprise a second graph node.

(289) 24. The method of 23 wherein the second graph node defines a sub-graph of the first sub-graph.

(290) 25. The method of 1 wherein the executable graph-based model comprises an edge node comprising connective relationships between at least two sub-graphs of the executable graph-based model.

(291) 26. The method of 1 further comprising, after the step of causing execution: obtaining, by the processing circuitry, one or more outcomes of said execution of processing logic within the overlay structure; and outputting, by the processing circuitry, the one or more outcomes.

(292) 27. The method of 1 wherein nodes within each sub-graph are semantically related.

(293) 28. The method of 1 wherein the executable graph-based model comprises a shared node within the hierarchical structure of each of the plurality of sub-graphs.

(294) 29. A system for dynamic execution of sub-graphs within executable graph-based models, the system comprising processing circuitry and a memory unit operatively coupled to the processing circuitry and having instructions stored thereon that, when executed by the processing circuitry, cause the processing circuitry to: obtain an executable graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein each sub-graph defines a hierarchical structure of related nodes; receive a stimulus and a context associated with the stimulus; in response to the stimulus being received, map, based on the context, the stimulus to a first sub-graph of the executable graph-based model;

and cause execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with one or more nodes of the first sub-graph.

(295) 30. A non-transitory computer-readable medium for dynamic execution of sub-graphs within executable graph-based models, the non-transitory computer-readable medium storing instructions which, when executed by processing circuitry, cause the processing circuitry to: obtain an executable graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein each sub-graph defines a hierarchical structure of related nodes; receive a stimulus and a context associated with the stimulus; in response to the stimulus being received, map, based on the context, the stimulus to a first sub-graph of the executable graph-based model; and cause execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with one or more nodes of the first sub-graph.

(296) 31. A device comprising: a display; an interface unit configured to receive one or more user inputs; and processing circuitry operatively coupled to the display and the interface unit, wherein the processing circuitry is configured to: obtain an executable graph-based model comprising a first sub-graph and at least one node, wherein the first sub-graph comprises a hierarchical structure of one or more related nodes; display, on the display of the device, a graphical representation of the executable graph-based model, the graphical representation comprising a first display element associated with the first sub-graph and at least one second display element associated with the at least one node, wherein the first display element represents the first sub-graph at a first level of detail; display, on the display of the device, a first selector associated with the first display element; receive, from the interface unit, a first user input associated with the first selector; and in response to the first user input being received, replace the first display element of the graphical representation of the executable graph-based model with a third display element whilst maintaining display of the at least one second display element, wherein the third display element represents the first sub-graph at a second level of detail.

(297) 32. The device of 31 wherein the second level of detail is greater than the first level of detail.

(298) 33. The device of 32 wherein the first display element comprises a first shape.

(299) 34. The device of 33 wherein the third display element comprises the first shape and one or more shapes associated with the one or more related nodes of the hierarchical structure of the first sub-graph.

(300) 35. The device of 34 wherein the one or more shapes are positioned circumferentially around the first shape.

(301) 36. The device of 34 wherein the one or more shapes are positioned within the first shape.

(302) 37. The device of 31 wherein the second level of detail is less than the first level of detail.

(303) 38. The device of 31 wherein the graphical representation of the executable graph-based model further comprises one or more further graphical representations of one or more further elements of the executable graph-based model.

(304) 39. The device of 38 wherein the first display element is replaced by the third display element, in response to the first user input being received, whilst maintaining display of the one or more further graphical representations.

(305) 40. The device of 38 wherein the one or more further elements of the executable graph-based model include one or more of: a node, an edge, a role, and an overlay.

(306) 41. The device of 31 wherein the first selector is displayed proximate the first display element.

(307) 42. The device of 31 wherein the first selector and the first display element are contiguous.

(308) 43. A method for interactive visualization of sub-graphs within an executable graph-based model, the method comprising: obtaining, by processing circuitry, an executable graph-based model comprising a first sub-graph and at least one node, wherein the first sub-graph comprises a hierarchical structure of one or more related nodes; displaying, by the processing circuitry, a

graphical representation of the executable graph-based model, the graphical representation comprising a first display element associated with the first sub-graph and at least one second display element associated with the at least one node, wherein the first display element represents the first sub-graph at a first level of detail; displaying, by the processing circuitry, a first selector associated with the first display element; receiving, by the processing circuitry, a first user input associated with the first selector; and in response to the first user input being received, replacing, by the processing circuitry, the first display element of the graphical representation of the executable graph-based model with a third display element whilst maintaining display of the at least one second display element, wherein the third display element represents the first sub-graph at a second level of detail.

(309) 44. A non-transitory computer-readable medium for interactive visualization of sub-graphs within an executable graph-based model, the non-transitory computer-readable medium storing instructions which, when executed by processing circuitry, cause the processing circuitry to: obtain an executable graph-based model comprising a first sub-graph and at least one node, wherein the first sub-graph comprises a hierarchical structure of one or more related nodes; display a graphical representation of the executable graph-based model, the graphical representation comprising a first display element associated with the first sub-graph and at least one second display element associated with the at least one node, wherein the first display element represents the first sub-graph at a first level of detail; display a first selector associated with the first display element; receive a first user input associated with the first selector; and in response to the first user input being received, replace the first display element of the graphical representation of the executable graph-based model with a third display element whilst maintaining display of the at least one second display element, wherein the third display element represents the first sub-graph at a second level of detail.

(310) 45. A method for loading of sub-graphs in executable graph-based models, the method comprising: obtaining, by processing circuitry and based on a first identifier, a first manifest state, wherein the first manifest state comprises the first identifier and a second identifier; obtaining, by the processing circuitry and based on the second identifier, a second manifest state, wherein the second manifest state comprises the second identifier; generating, by the processing circuitry, a first manifest from the first manifest state, wherein the first manifest comprises a first node state storage location, the first identifier, and the second identifier; generating, by the processing circuitry, a second manifest from the second manifest state, wherein the second manifest comprises a second node state storage location and the second identifier; obtaining, by the processing circuitry, a first node state and a second node state from the first node state storage location and the second node state storage location respectively; generating, by the processing circuitry, a first node based on the first manifest and the first node state, wherein the first node comprises the first identifier and the second identifier; generating, by the processing circuitry, a first graph node based on the second manifest and the second node state, wherein the first graph node comprises the second identifier; and generating, by the processing circuitry, a first sub-graph of an executable graph-based model by associating the first node with the first graph node.

(311) 46. The method of 45 wherein the first node is a second graph node defining a second sub-graph within the executable graph-based model.

(312) 47. The method of 46 wherein the second sub-graph is a child of the first sub-graph.

(313) 48. The method of 45 wherein the second manifest state further comprises a third identifier associated with a shared sub-graph.

(314) 49. The method of 48 further comprising: loading, by the processing circuitry, the shared sub-graph-based on the third identifier; and associating, by the processing circuitry, the shared sub-graph with the first sub-graph.

(315) 50. A device comprising processing circuitry and a memory storing instructions which, when executed by the processing circuitry, cause the processing circuitry to: obtain, based on a first

identifier, a first manifest state, wherein the first manifest state comprises the first identifier and a second identifier; obtain, based on the second identifier, a second manifest state, wherein the second manifest state comprises the second identifier; generate a first manifest from the first manifest state, wherein the first manifest comprises a first node state storage location, the first identifier, and the second identifier; generate a second manifest from the second manifest state, wherein the second manifest comprises a second node state storage location and the second identifier; obtain a first node state and a second node state from the first node state storage location and the second node state storage location respectively; generate a first node based on the first manifest and the first node state, wherein the first node comprises the first identifier and the second identifier; generate a first graph node based on the second manifest and the second node state, wherein the first graph node comprises the second identifier; and generate a first sub-graph of an executable graph-based model by associating the first node with the first graph node.

(316) 51. A non-transitory computer-readable medium storing instructions which, when executed by processing circuitry, cause the processing circuitry to: obtain, based on a first identifier, a first manifest state, wherein the first manifest state comprises the first identifier and a second identifier; obtain, based on the second identifier, a second manifest state, wherein the second manifest state comprises the second identifier; generate a first manifest from the first manifest state, wherein the first manifest comprises a first node state storage location, the first identifier, and the second identifier; generate a second manifest from the second manifest state, wherein the second manifest comprises a second node state storage location and the second identifier; obtain a first node state and a second node state from the first node state storage location and the second node state storage location respectively; generate a first node based on the first manifest and the first node state, wherein the first node comprises the first identifier and the second identifier; generate a first graph node based on the second manifest and the second node state, wherein the first graph node comprises the second identifier; and generate a first sub-graph of an executable graph-based model by associating the first node with the first graph node.

(317) 52. A method for storage management of sub-graphs in executable graph-based models, the method comprising: obtaining, by processing circuitry, a first sub-graph of an executable graph-based model, the first sub-graph comprising a first graph node and a second node, wherein the first graph node defines the first sub-graph and is associated with a first identifier; extracting, by the processing circuitry, a first node state from first graph node, wherein the first node state comprises the first identifier; extracting, by the processing circuitry, a second node state from the second node, wherein the second node state comprises the first identifier and a second identifier associated with the second node; determining, by the processing circuitry, a storage location for the node state of each of the first graph node and the second node; generating, by the processing circuitry, a first sub-graph manifest associated with the first graph node, wherein the first sub-graph manifest comprises the first identifier and the storage location for the first node state; generating, by the processing circuitry, a first node manifest associated with the second node, wherein the first node manifest comprises the first identifier, the second identifier, and the storage location for the second node state; generating, by the processing circuitry, a first manifest state for the first sub-graph manifest, wherein the first manifest state comprises the first identifier; generating, by the processing circuitry, a second manifest state for the first node manifest, wherein the second manifest state comprises the first identifier and the second identifier; and storing, by the processing circuitry, the first manifest state and the second manifest state.

(318) 53. A device comprising processing circuitry and a memory storing instructions which, when executed by the processing circuitry, cause the processing circuitry to: obtain a first sub-graph of an executable graph-based model, the first sub-graph comprising a first graph node and a second node, wherein the first graph node defines the first sub-graph and is associated with a first identifier; extract a first node state from first graph node, wherein the first node state comprises the first identifier; extract a second node state from the second node, wherein the second node state

comprises the first identifier and a second identifier associated with the second node; determine a storage location for the node state of each of the first graph node and the second node; generate a first sub-graph manifest associated with the first graph node, wherein the first sub-graph manifest comprises the first identifier and the storage location for the first node state; generate a first node manifest associated with the second node, wherein the first node manifest comprises the first identifier, the second identifier, and the storage location for the second node state; generate a first manifest state for the first sub-graph manifest, wherein the first manifest state comprises the first identifier; generate a second manifest state for the first node manifest, wherein the second manifest state comprises the first identifier and the second identifier; and store the first manifest state and the second manifest state.

(319) 54. A non-transitory computer-readable medium storing instructions which, when executed by processing circuitry, cause the processing circuitry to: obtain a first sub-graph of an executable graph-based model, the first sub-graph comprising a first graph node and a second node, wherein the first graph node defines the first sub-graph and is associated with a first identifier; extract a first node state from first graph node, wherein the first node state comprises the first identifier; extract a second node state from the second node, wherein the second node state comprises the first identifier and a second identifier associated with the second node; determine a storage location for the node state of each of the first graph node and the second node; generate a first sub-graph manifest associated with the first graph node, wherein the first sub-graph manifest comprises the first identifier and the storage location for the first node state; generate a first node manifest associated with the second node, wherein the first node manifest comprises the first identifier, the second identifier, and the storage location for the second node state; generate a first manifest state for the first sub-graph manifest, wherein the first manifest state comprises the first identifier; generate a second manifest state for the first node manifest, wherein the second manifest state comprises the first identifier and the second identifier; and store the first manifest state and the second manifest state.

Claims

1. A method for dynamic execution of sub-graphs within executable graph-based models, the method comprising: obtaining, by processing circuitry, a graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein: each sub-graph defines a hierarchical structure of related nodes, each sub-graph is defined by a graph node within the graph-based model, and the graph node that defines each sub-graph encapsulates one or more nodes included within the sub-graph; generating, by the processing circuitry, during a run-time of the processing circuitry, an executable graph-based model based on the obtained graph-based model; receiving, by the processing circuitry, a stimulus and a context associated with the stimulus; in response to the stimulus being received, mapping, by the processing circuitry and based on the context, the stimulus to a first sub-graph of the plurality of sub-graphs of the executable graph-based model; and causing, by the processing circuitry, execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with the one or more nodes of the first sub-graph.

2. The method of claim 1, wherein the context is associated with a target executable node of the first sub-graph, the target executable node comprising a composition of a first node of the first sub-graph and a first overlay of the overlay structure, the first overlay comprising processing logic operable to interact with the first node such that the step of causing execution of processing logic within the overlay structure comprises causing execution of said processing logic of the first overlay.

3. The method of claim 2, wherein the graph node that defines each sub-graph is one of a group consisting of: a leaf graph node and a composite graph node, wherein the first sub-graph of the

executable graph-based model is defined by a first graph node, and wherein the first graph node is the composite graph node and is a root node of the first sub-graph.

4. The method of claim 3, further comprising, in response to the stimulus being received and prior to the step of causing execution: traversing, by the processing circuitry, from the target executable node to the root node of the first sub-graph thereby determining a traversal path comprising a first set of nodes from the root node to the target executable node, wherein the first set of nodes includes the target executable node.

5. The method of claim 4, wherein the step of causing execution comprises: identifying, by the processing circuitry, a plurality of executable nodes within the first set of nodes, each of the plurality of executable nodes comprising a composition of a respective node of the first sub-graph and at least one overlay of the overlay structure; and causing, by the processing circuitry, execution of processing logic associated with each overlay of the plurality of executable nodes.

6. The method of claim 4 further comprising, after the step of traversing and prior to the step of causing execution: loading, by the processing circuitry and to memory operatively coupled to the processing circuitry, the first set of nodes of the traversal path, wherein the first set of nodes are loaded in order from the root node to the target executable node.

7. The method of claim 6, wherein the first set of nodes comprises a node associated with at least one overlay within the overlay structure such that the step of loading the first set of nodes comprises: generating, by the processing circuitry, an executable node comprising a composition of the node and the at least one overlay.

8. The method of claim 6, wherein the first graph node is associated with at least one overlay within the overlay structure such that the first graph node is an executable node.

9. The method of claim 8, wherein the step of loading the first set of nodes comprises: generating, by the processing circuitry, a first set of executable nodes based on the first set of nodes and the at least one overlay, wherein each executable node within the first set of executable nodes comprises a composition of a respective node of the first set of nodes and the at least one overlay.

10. The method of claim 6, further comprising, after the step of causing execution: unloading, by the processing circuitry and from memory operatively coupled to the processing circuitry, one or more nodes within the first set of nodes, wherein the one or more nodes are unloaded in order from the target executable node to the root node.

11. The method of claim 10, wherein the one or more nodes are unloaded when an unloading criterion is met.

12. The method of claim 11, wherein the unloading criterion is met after a predetermined period of time has elapsed from the step of causing execution.

13. The method of claim 11, wherein the unloading criterion is based on the stimulus.

14. The method of claim 3, wherein the first graph node encapsulates one or more further nodes within the first sub-graph.

15. The method of claim 14, wherein the one or more further nodes comprise one or more of: a data node, a value node, and a role node.

16. The method of claim 14, wherein the one or more further nodes comprise an edge node indicative of a first connective relationship between the first graph node and the edge node.

17. The method of claim 16, wherein the first connective relationship is defined by a first role.

18. The method of claim 16, wherein the edge node further comprises a second connective relationship between the edge node and a second node.

19. The method of claim 18, wherein the first sub-graph comprises the second node.

20. The method of claim 18, wherein the second node is not a part of the first sub-graph.

21. The method of claim 18, wherein the second node is a graph node which defines a sub-graph of the executable graph-based model.

22. The method of claim 18, wherein the second connective relationship is defined by a second role.

23. The method of claim 14, wherein the one or more further nodes comprise a second graph node.
24. The method of claim 23, wherein the second graph node defines a sub-graph of the first sub-graph.
25. The method of claim 1, wherein the executable graph-based model comprises an edge node indicative of connective relationships between at least two sub-graphs of the executable graph-based model.
26. The method of claim 1, further comprising, after the step of causing execution: obtaining, by the processing circuitry, one or more outcomes of said execution of processing logic within the overlay structure; and outputting, by the processing circuitry, the one or more outcomes.
27. The method of claim 1, wherein nodes within each sub-graph are semantically related.
28. The method of claim 1, wherein the executable graph-based model comprises a shared node within the hierarchical structure of each of the plurality of sub-graphs.
29. A system for dynamic execution of sub-graphs within executable graph-based models, the system comprising processing circuitry and a memory unit operatively coupled to the processing circuitry and having instructions stored thereon that, when executed by the processing circuitry, cause the processing circuitry to: obtain a graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein: each sub-graph defines a hierarchical structure of related nodes, each sub-graph is defined by a graph node in the graph-based model, and the graph node that defines each sub-graph encapsulates one or more nodes included within the sub-graph; generate, during a run-time of the processing circuitry, an executable graph-based model based on the obtained graph-based model; receive a stimulus and a context associated with the stimulus; in response to the stimulus being received, map, based on the context, the stimulus to a first sub-graph of the plurality of sub-graphs of the executable graph-based model; and cause execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with the one or more nodes of the first sub-graph.
30. A non-transitory computer-readable medium for dynamic execution of sub-graphs within executable graph-based models, the non-transitory computer-readable medium storing instructions which, when executed by processing circuitry, cause the processing circuitry to: obtain a graph-based model comprising a plurality of sub-graphs and an overlay structure comprising processing logic associated with the plurality of sub-graphs, wherein: each sub-graph defines a hierarchical structure of related nodes, each sub-graph is defined by a graph node, and the graph node that defines each sub-graph encapsulates one or more nodes included within the sub-graph; generating during a run-time of the processing circuitry, an executable graph-based model based on the obtained graph-based model; receive a stimulus and a context associated with the stimulus; in response to the stimulus being received, map, based on the context, the stimulus to a first sub-graph of the executable graph-based model; and cause execution of processing logic within the overlay structure based on the mapping, wherein said processing logic is associated with the one or more nodes of the first sub-graph.
-