

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12387089
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Shao; Yakun et al.

Efficient neural network accelerator dataflows

Abstract

A distributed deep neural net (DNN) utilizing a distributed, tile-based architecture includes multiple chips, each with a central processing element, a global memory buffer, and a plurality of additional processing elements. Each additional processing element includes a weight buffer, an activation buffer, and vector multiply-accumulate units to combine, in parallel, the weight values and the activation values using stationary data flows.

Inventors: Shao; Yakun (Santa Clara, CA), Venkatesan; Rangharajan (Sunnyvale, CA), Wang; Miaorong (Cambridge, MA), Smith; Daniel (Los Gatos, CA), Dally; William James (Incline Village, NV), Emer; Joel (Acton, MA), Keckler; Stephen W. (Austin, TX), Khailany; Brucek (Austin, TX)

Applicant: NVIDIA Corp. (Santa Clara, CA)

Family ID: 72423013

Assignee: Nvidia Corp. (Santa Clara, CA)

Appl. No.: 17/530852

Filed: November 19, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220076110 A1	Mar. 10, 2022

Related U.S. Application Data

division parent-doc US 16672918 20191104 US 11270197 child-doc US 17530852
us-provisional-application US 62817413 20190312

Publication Classification

Int. Cl.: G06N3/063 (20230101); G06F9/30 (20180101); G06F9/38 (20180101); G06F17/16 (20060101); G06N3/08 (20230101)

U.S. Cl.:

CPC G06N3/063 (20130101); G06F9/30036 (20130101); G06F9/3877 (20130101); G06F17/16 (20130101); G06N3/08 (20130101);

Field of Classification Search

CPC: G06N (3/063); G06N (3/08); G06N (3/045); G06F (9/30036); G06F (9/3877); G06F (17/16); G06F (9/3001)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5008833	12/1990	Agranat et al.	N/A	N/A
5091864	12/1991	Baji et al.	N/A	N/A
5297232	12/1993	Murphy	N/A	N/A
5319587	12/1993	White	N/A	N/A
5704016	12/1996	Shigematsu et al.	N/A	N/A
5768476	12/1997	Sugaya et al.	N/A	N/A
6560582	12/2002	Woodall	N/A	N/A
7082419	12/2005	Lightowler	N/A	N/A
10817260	12/2019	Huang et al.	N/A	N/A
10929746	12/2020	Litvak et al.	N/A	N/A
2003/0061184	12/2002	Masgonty et al.	N/A	N/A
2009/0210218	12/2008	Collobert et al.	N/A	N/A
2018/0004689	12/2017	Marfatia et al.	N/A	N/A
2018/0046897	12/2017	Kang et al.	N/A	N/A
2018/0046916	12/2017	Dally	N/A	G06F 7/523
2018/0121796	12/2017	Deisher et al.	N/A	N/A
2018/0173676	12/2017	Tsai et al.	N/A	N/A
2019/0050717	12/2018	Temam et al.	N/A	N/A
2019/0138891	12/2018	Kim et al.	N/A	N/A
2019/0179975	12/2018	Rassekh et al.	N/A	N/A
2019/0180170	12/2018	Huang et al.	N/A	N/A
2019/0340510	12/2018	Li et al.	N/A	N/A
2019/0370645	12/2018	Lee et al.	N/A	N/A
2019/0392287	12/2018	Ovsiannikov et al.	N/A	N/A
2021/0004668	12/2020	Moshovos et al.	N/A	N/A
2021/0125046	12/2020	Moshovos et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
107463932	12/2016	CN	N/A

OTHER PUBLICATIONS

Yavits, Leonid, Amir Morad, and Ran Ginosar. "Sparse matrix multiplication on an associative processor." IEEE Transactions on Parallel and Distributed Systems 26.11 (2014): 3175-3183. (Year: 2014). cited by applicant

Angshuman Parashar et al. "SCN N: An Accelerator for Compressed-sparse Convolutional Neural Networks", (Year: 2017). cited by applicant

Duckhwan Kim et al. "Neurocube: A Programmable Digital Neuromorphic Architecture with High-Density 3D Memory", (Year: 2016). cited by applicant

Eunhyeok Park et al. "Energy-efficient Neural Network Accelerator Based on Outlier-aware Low-precision Computation", Jun. 1, 2018 (Year: 2018). cited by applicant

U.S. Appl. No. 16/517,431, filed Jul. 19, 2019, Yakun Shao. cited by applicant

Andri, YodaN N: An Ultra-Low Power Convolutional Neural Network Accelerator Based on Binary Weights, IEEE, pp. 238-241 (Year: 2016). cited by applicant

Han, EIE: Efficient Inference Engine on Compressed Deep Neural Network, IEEE, pp. 243-254 (Year: 2016). cited by applicant

Kartik Hegde et al. "UCNN: Exploiting Computational Reuse in Deep Neural Networks via Weight Repetition", Apr. 18, 2018. (Year: 2018). cited by applicant

Streat, Lennard, Dhiressha Kudithipudi, and Kevin Gomez. "Non-volatile hierarchical temporal memory: Hardware for spatial pooling." arXiv preprint arXiv:1611.02792 (2016). (Year: 2016). cited by applicant

F. Sijstermans. The NVIDIA deep learning accelerator. In Hot Chips, 2018. cited by applicant

K. Kwon et al. Co-design of deep neural nets and neural net accelerators for embedded vision applications. In Proc. DAC, 2018. cited by applicant

N. P. Jouppi et al. In-datacenter performance analysis of a tensor processing unit. In Proc. ISCA, 2017. cited by applicant

T. Chen et al. DianNao: A small-footprint high-throughput accelerator for ubiquitous machine-learning. In Proc. ASPLOS, 2014. cited by applicant

W. Lu et al. FlexFlow: A flexible dataflow accelerator architecture for convolutional neural networks. In Proc. HPCA, 2017. cited by applicant

Y. Chen, J. Emer, and V. Sze. Eyeriss: A spatial architecture for energy-efficient dataflow for convolutional neural networks. In Proc. ISCA, 2016. cited by applicant

Primary Examiner: Roy; Sanchita

Attorney, Agent or Firm: Carnelian Law

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application claims priority and benefit under 35 U.S.C. 119 (e) to U.S. Application Ser. No. 62/817,413, filed on Mar. 12, 2019, the contents of which are incorporated herein by reference in their entirety.

BACKGROUND

(1) Deep neural networks (DNNs) are a class of neural network that has emerged as a key approach for solving complex problems across various technical fields. "Neural network" refers to an algorithm or computational system based on a collection of connected units or nodes called artificial neurons, which loosely model the neurons in a biological system. Each connection between neurons, like the synapses in a biological brain, can transmit a signal (an activation) from

one artificial neuron to another. An artificial neuron that receives a signal (the input activation) can process it and then signal additional artificial neurons (the output activation) connected to it. Applications of DNNs have diverse performance, accuracy, and power requirements depending on the implementation. Building dedicated DNNs for the requirements of particular implementations may be cost prohibitive due to high design complexity and manufacturing challenges.

Description

BRIEF DESCRIPTION OF THE SEVERAL VIEWS OF THE DRAWINGS

- (1) To easily identify the discussion of any particular element or act, the most significant digit or digits in a reference number refer to the figure number in which that element is first introduced.
- (2) FIG. 1 depicts a scalable neural network processor **100** in accordance with one embodiment.
- (3) FIG. 2 depicts a multi-die package **202** in accordance with one embodiment.
- (4) FIG. 3 depicts a neural network processor **300** implemented on a single chip in accordance with one embodiment.
- (5) FIG. 4 depicts a local processing element **400** in accordance with one embodiment.
- (6) FIG. 5 depicts a local processing element **500** in more detail, in accordance with one embodiment.
- (7) FIG. 6 depicts a global processing element **602** in accordance with one embodiment.
- (8) FIG. 7 depicts convolution **700** in accordance with one embodiment.
- (9) FIG. 8 depicts a chip-level tiling **800** in accordance with one embodiment.
- (10) FIG. 9 depicts a processing element-level tiling **900** in accordance with one embodiment.
- (11) FIG. 10 depicts neural network processing **1000** in accordance with one embodiment.
- (12) FIG. 11 illustrates a deeply-nested loop data flow **1100** in accordance with one embodiment.
- (13) FIG. 12 depicts a weight-stationary data flow **1200** in accordance with one embodiment.
- (14) FIG. 13 depicts an output-stationary data flow **1300** in accordance with one embodiment.
- (15) FIG. 14 depicts an input-stationary data flow **1400** in accordance with one embodiment.
- (16) FIG. 15 depicts a weight-stationary-local-output-stationary data flow **1500** in accordance with one embodiment.
- (17) FIG. 16 depicts an output-stationary-local-weight-stationary data flow **1600** in accordance with one embodiment.
- (18) FIG. 17 depicts a parallel processing unit **1700** in accordance with one embodiment.
- (19) FIG. 18 depicts a general processing cluster **1800** in accordance with one embodiment.
- (20) FIG. 19 depicts a memory partition unit **1900** in accordance with one embodiment.
- (21) FIG. 20 depicts a streaming multiprocessor **2000** in accordance with one embodiment.
- (22) FIG. 21 depicts a processing system **2100** in accordance with one embodiment.
- (23) FIG. 22 depicts an exemplary processing system **2200** in accordance with another embodiment.

DETAILED DESCRIPTION

- (24) Disclosed herein is an efficient, tiled DNN processor that utilizes a scalable design. The disclosed system comprises beneficial features including: 1) a fully distributed, tile-based architecture, 2) flexible and efficient weight and activation tiling at the processing element (PE) level, chip-level, and in some embodiments, package-level, improving data locality and reducing communication cost, and 3) multi-level dataflows, improving the data reuse and energy efficiency.
- (25) The architecture utilizes a data path designed to account for the low computation-to-memory ratio of neural network layers. The data path includes, in some implementations, both local and global processing elements. Each local processing element comprises logic to perform localized multiply-accumulation of weights and input activations, and post-processing such as ReLu, MaxPool, etc. Memory buffers in the form of collectors and register files may be disposed into the

data path within and/or between processing elements.

(26) The architecture provides multi-level memory and computation hierarchy that exploits both weight and output activation locality to improve the energy efficiency of neural network execution. Conventional neural network accelerator designs only leverage the reuse opportunity of the innermost execution level (e.g., loop), whereas the disclosed architecture provides a multi-level memory and processing hierarchy to exploit data reuse opportunities across multiple loop levels, thus enabling a diverse set of energy-efficient data flows. For example, instead of capturing temporal reuse only for weights or outputs, multi-level dataflows may be implemented that exploit both weight and partial sum reuse during the execution.

(27) To efficiently implement a particular data flow, each local processing element may utilize one or more collectors (e.g., small register-files): one in front of a weight buffer, another one in front of an accumulation buffer, and another in front of an input activation buffer. These collectors filter out (reduce) expensive reads and writes to the weight and partial sum buffers (e.g., SRAMs), leading to an overall energy efficiency improvement. The global processing elements and/or the chip may provide additional storage (e.g., a global or shared register file) and processing capability in the data path of neural network computations.

(28) The disclosed architecture provides a heterogeneous-tile-based computational platform for different types neural network calculations. In addition to dense convolution, many neural networks perform element-wise calculation and depth-wise convolution. To facilitate such computations the architecture includes two general types of processing element. The first type, called local processing elements, specialize in executing dense convolution with significant data reuse. The second type, called global processing elements, provide second-level storage for the local processing elements during dense convolution. In addition, the global processing elements may perform element-wise operations and depth-wise convolution at a low compute-to-memory ratio without communicating large amounts of data through layers of the neural network.

(29) FIG. 1 illustrates a scalable neural network processor **100** in one embodiment. Weights for a deep neural network are tiled spatially across on-die memories (storage **104**) via an interconnect **102**. “Weights” refers to values with which activations are multiplied to increase or decrease the impact of the activation values in an activation function. Herein the terms “die” and “chip” are used interchangeably. Each chip/die comprises of one or more (typically multiple) processing elements **106**. The weights may be partitioned among any number of chips (dice), for example 1-36 dice. As depicted, FIG. 1 comprises 36 dice. The scalable neural network processor **100** may be implemented by multiple dice in a single package/device/printed circuit board, or across multiple packages/devices/printed circuit boards.

(30) Deep neural network applications can differ significantly in their requirements. For example, typical data center inference applications such as image recognition may prioritize performance and scalability at low latency and may be willing to sacrifice classification accuracy, while inference for autonomous driving workloads may prioritize energy efficiency within real-time constraints while maintaining the best achievable network accuracy. Because application-specific inference accelerators can provide significant performance and power advantages compared to general-purpose solutions, it is desirable to enable custom solutions on a common architecture for different target markets and applications.

(31) A multi-die package **202** embodiment for implementing a DNN accelerator is depicted in FIG. 2. The multi-die package **202** may be a semiconductor package comprising a plurality of dice **204** (chips). Each of the dice **204** comprises a plurality of processing elements **206**, a global buffer **208**, and a controller **210** (e.g., an open-source RISC-V processor). The elements of each chip/die communicate via a network-on-a-chip router **212**. The multiple chips in a package communicate with one another via a network-on-a-package router **214**, and may also communicate with a host **216** system comprising DRAM **220** or other memory, via a Field Programmable Gate Array (FPGA **218**), Joint Test Action Group (JTAG) logic, or other interface technology as known in the art.

(32) Some or all of the processing elements are local processing elements comprising a weight buffer to receive and store weight values for a deep neural network. “Weight buffer” refers to a buffer storing weight values. The local processing elements comprise an activation buffer to receive activation values for the deep neural network. “Activation buffer” refers to a memory buffer utilized to store activation values (activations) utilized in a neural network computation.

Activations are computed by each neuron in a neural network layer using an activation function, also sometimes called a ‘transfer function’. Activations may be simple binary values (e.g., “1” or “0” representing “ON” or “OFF”) or they may take on a range of values for some activation functions. The weight buffer and activation buffer may be separate elements within each processing element. “Activations” refers to the output values of neurons in a neural network, computed based at least in part on weights input to the neuron and an activation function of the neuron. Activations are also called ‘activation values’. “Buffer” refers to a memory storing values that are inputs to or results from a calculation. The local processing elements further comprise a plurality of multiply-accumulate units to combine, in parallel, the weight values and the activation values, to generate partial sums.

(33) The multi-die package **202** may be configured to distribute the weight values and the activation values among the local processing elements spatially and temporally (over time). The global memory buffer of each chip may act as a second-level buffer for the activation values during computation. “Second-level buffer” refers a memory where values are stored and retrieved from when the values are needed for computation but are not available in the first-level buffer. Herein, the chip global buffer may act as a second-level buffer to the first-level activation buffers of the chip's processing elements. “Global memory buffer” refers to a buffer available for utilization by all or at least a plurality of processing elements on a chip. The distribution of weights and activations during computation may be carried out by the chip's controller **210**. The controller **210** or local controllers of any of the processing elements **206** may be configured by instructions stored in a memory to carry out various data flows described below. A memory configured in such a manner may conveniently be referred to herein as “logic”. The location of such logic is a design choice. The memory storing such instructions may be any of the memories depicted in the figures, or a different memory not depicted.

(34) FIG. 3 depicts a neural network processor **300** embodied on a single chip. The neural network processor **300** may utilize a fixed point data path between a plurality of processing elements **206**. The neural network processor **300** also comprises the aforementioned global buffer **208** and controller **210**, which for example may be a RISC-V processor. The processing elements **206** and global buffer **208** communicate via the network-on-a-chip router **212** or other interconnect technology (see the GPU implementations, described further below). If a router is utilized, it may be implemented centrally or in distributed fashion as routers on each of the processing elements **206**. The processing elements **206** utilize the router/interconnect to communicate with processing elements on the same package, or in some embodiments across packages via a network-on-a-package router **214**.

(35) FIG. 4 depicts, and at a high level, an exemplary local processing element **400**. The processing element **400** includes a plurality of vector multiply-accumulate units **402**, a weight buffer **404**, an activation buffer **406**, a router **408**, a controller **410**, an accumulation memory buffer **412**, and a post-processor **414**. The activation buffer **406** may, in one embodiment, be implemented as a dual-ported SRAM to receive activation values from the global buffer **208** or from other local or global processing elements, via the router **408** or other interconnect. The router **408** may be a component of a distributed network-on-a-chip router **212** that in one embodiment comprises a serializer/deserializer, packetizer, arbitrator, Advanced extensible Interface, and other components known in the art.

(36) The weight buffer **404** may, in one embodiment, be implemented as a single-ported SRAM storing weight values. The weight values used by the vector multiply-accumulate units **402** may be

“weight-stationary”, meaning they are not updated each clock cycle, but instead are updated only after the output activation values are computed for a particular layer of the deep neural network. (37) The accumulation memory buffer **412** may comprise one or more SRAM devices to store the output activations computed by the vector multiply-accumulate units **402**. The router **408** communicates these output activations and control signals from the processing element **400** to other processing elements. “Output activation” refers to an activation output by a neuron in a neural network. An output activation is typically computed based on the input activations to the neuron and the weights applied to the input activations. “Input activation” refers to an activation received by a neuron in a neural network.

(38) The processing element **400** may perform all operations of convolutional and fully-connected layers of a DNN efficiently, including multiply-accumulate, truncation, scaling, bias addition, ReLU, and pooling (these last five in the post-processor **414**). The vector multiply-accumulate units **402** may operate on the same inputs using different filters. In one embodiment, each of the vector multiply-accumulate units **402** performs an eight-input-channel dot product and accumulates the result into the accumulation memory buffer **412** on each clock cycle. The weights stored in the weight buffer **404** are unchanged until the entire computation of output activations completes. Each processing element **400** reads the input activations in the activation buffer **406**, performs the multiply-accumulate operations, and writes output activations to the accumulation memory buffer **412** on every clock cycle. The frequency at which the weight buffer **404** is accessed depends on the input activation matrix dimensions and the number of filters utilized.

(39) The vector multiply-accumulate units **402** of each processing element **400** computes a portion of a wide dot-product-accumulate as a partial result and forwards the partial result to neighboring processing elements. “Dot-product-accumulate” refers to the computation of a dot product. A dot product is the sum of the products of the corresponding entries of the two sequences (vectors) of numbers. Dot products are efficiently computed using vector multiply-accumulate units. “Multiply-accumulate unit” refers to a data processing circuit that carries out multiply-accumulate operations, which involve computing the product of two numbers and adding that product to an accumulator. Multiply-accumulate units may be referred to herein by their acronym, MAC or MAC unit. A multiply-accumulate unit carries out computations of the form $a \leftarrow a + (b * c)$. A vector multiply-accumulate unit computes the product of two vectors using an array of multipliers, then performs a reduction operation by adding all the outputs of multipliers to produce a partial sum, which is then added to an accumulator.

(40) The partial results are transformed into a final result by the post-processor **414** and communicated to the global buffer **208**. The global buffer **208** acts as a staging area for the final multiply-accumulate results between layers of the deep neural network.

(41) The accumulation memory buffer **412** receives outputs from the vector multiply-accumulate units **402**. The central controller **210** distributes the weight values and activation values among the processing elements and utilizes the global memory buffer as a second-level buffer for the activation values. When processing images, the controller **210** configures processing by layers of the deep neural network spatially across the processing elements by input/output channel dimensions and temporally by image height/width.

(42) The global buffer **208** stores both input activations and output activations from the processing elements **206** for distribution by the aforementioned transceivers to the processing elements via multicast. “Multicast” refers to a group communication mechanism whereby transmission of data is addressed to a group of destination devices (e.g., processing elements) simultaneously. Multicast can implement one-to-many or many-to-many distribution. Each of the processing elements **206** includes a router **408** to communicate, in one embodiment, 64 bits of data in, and 64 bits of data out, per clock cycle. This enables accumulation of partial sums for wide dot products that have their computation spatially tiled across the processing elements **206**.

(43) FIG. 5 depicts an exemplary local processing element **500** in more detail. The processing

element **500** includes the aforementioned vector multiply-accumulate units **402**, weight buffer **404**, activation buffer **406**, router **408**, controller **410**, accumulation memory buffer **412**, and post-processor **414** (e.g., the post-processor **414**). Also depicted are a weight collector **522** interposed between the weight buffer **404** and the vector multiply-accumulate units **402**, and an accumulation collector **524** interposed between the vector multiply-accumulate units **402** and the accumulation memory buffer **412**. The accumulation collector **524** may also be referred to herein as an “output collector”. Also depicted are various memory buffer managers that may be utilized (e.g., weight memory buffer manager **512**, activation memory buffer manager **518**, and accumulation memory buffer manager **520**). “Memory buffer manager” refers to logic for managing the contents of a memory buffer, for example managing the availability of certain data (e.g., weights, activations) in the buffer when requested by a processing element.

(44) The processing element **500** includes vector multiply-accumulate units **402** of which a number N are operational for a given data flow. Each vector multiply accumulate unit **526** performs V multiplications and additions per clock cycle. Thus, in every clock cycle, the processing element **500** can multiply a weight matrix of dimensions $N \times V$ with an input activation vector of size V , to generate a partial-sum vector of size N . In other words, each of the vector multiply-accumulate units **402** can perform a V -wide dot product calculation per clock cycle. One or both N and V may be configurable at the controller **210**.

(45) The input activation buffer **406** has an operational size IA and the weight buffer **404** has an operational size W . One or both W and IA may be configurable at the controller **210**. The accumulation memory buffer **412** has an operational size of A . “Operational size” refers to a resource pool available for performing calculations during operation of a device, which may be less than the total or maximum size of the resource pool. The operational size may be configurable using registers or other settings (e.g., for higher performance or less power consumption).

(46) Each of the vector multiply-accumulate units **402** includes a weight collector **522** buffer having a configurable depth (e.g., number of distinct registers or addresses in a register file used by the vector multiply-accumulate units **402** during computations) of WD and a width $V \times WP$ (WP is also called the weight precision). Across the N vector multiply-accumulate units **402**, the total weight collector width is $V \times N \times WP$. The input activations have width IAP . Each of the vector multiply-accumulate units **402** also includes an accumulation collector **524** having a configurable operational depth AD and width AP (AP is also called the accumulator precision). The accumulation collector **524** has a total width of $N \times AP$ across the N vector multiply-accumulate units **402**. The V -wide dot products and N -sized partial-sum vector may thus be computed by each vector multiply accumulate unit **526** at mixed precision. Some or all of WD , WP , IAP , AD , and AP may be configurable by the controller **210**.

(47) The weight buffer **404** read (output) port is $WP \times N \times V$ bits wide and is able to supply different weight vectors to different ones of the vector multiply-accumulate units **402**. The activation buffer **406** is $IAP \times V$ bits wide because the same IA vector is provided in parallel to all N vector multiply-accumulate units **402**.

(48) The values of V and N may be adjusted for more or less parallelism and reuse of weights, for example. Based on the configuration of N and V , other parameters such as W , IA , and A may be adjusted to ensure the vector multiply-accumulate units **402** stay busy during convolution calculation.

(49) The weight buffer **404** and the activation buffer **406** each have an associated address generator (address generator **514** and address generator **516**, respectively) that generates an address every cycle. “Address generator” refers to logic that calculates address values in a memory for reading or writing data from the address. The ordering of operations carried out by the vector multiply-accumulate units **402** is controlled by these address generators, which are configurable to support temporal reuse of weights or results in the accumulation collector **524** across clock cycles for different types of data flows. The depth WD of the weight collector **522** may be configurable to

enable different amounts of temporal reuse of partial sum values, depending on the requirements of the data flow. Likewise, the depth AD of the accumulation collector **524** may be configurable to enable different amounts of temporal reuse of weight values, depending on the requirements of the data flow.

(50) The processing element **500** may further comprise an input collector **528** disposed between the activation buffer **406** and the vector multiply-accumulate units **402**. An operational depth IC of the input collector **528** may be configured to set different levels of input activation stationary data flows, as described further below.

(51) Each of the weight buffer **404** and activation buffer **406** also have a buffer manager (weight memory buffer manager **512** and activation memory buffer manager **518**, respectively) responsive to the controller **410** and determining the availability of data to the vector multiply-accumulate units **402**. The dimensions of the address generators and the granularity of data movement from the weight buffer **404** and activation buffer **406** to the vector multiply-accumulate units **402** may in some embodiments be configurable at the controller **210**.

(52) The accumulation memory buffer **412** stores partial sums from all N vector multiply-accumulate units **402** and may be optimized to perform read-modify-write operations every cycle. Partial sums from the N vector multiply-accumulate units **402** are packed into vectors of width $AP \times N$ and stored in the accumulation memory buffer **412**. From there, they can be sent either directly to another processing element for cross-processing element reduction or to the post-processor **414** to produce final output activations. The post-processor **414** supports not only scaling and quantization but also ReLU and pooling operations to enable layer fusion.

(53) Input weights **502** arrive over the router **408** and are stored in the weight buffer **404**. Input activations **504** also arrive over the router **408** and are stored in the activation buffer **406**. Computed output activations **506** (after post-processing by the post-processor **414**) or partial sums **508** from the accumulation memory buffer **412** are output to the global buffer **208** or neighboring processing elements, respectively, via the router **408**. Cross-processing-element reductions **510** from said neighboring processing elements may be received by the router **408** and are accumulated in the accumulation memory buffer **412**.

(54) FIG. **6** illustrates the global processing element **602** in one embodiment. The global processing element **602** comprises a global memory buffer **604** with arbitrated memory banks **606** (e.g., a “scratchpad”), a controller **608** to carry out calculations on data in the arbitrated memory banks **606**, and an activation address generator **610** and a destination address generator **612** to generate source and destination addresses, respectively, for calculations. The global processing element **602** communicates with other processing elements via the router **408**.

(55) The data path **616** to and from the global memory buffer **604** comprises a register file **614** which may operate as a collector for one or more of input activations **618**, output activations **620**, and partial sums **622** to and from local processing elements, according to the requirements of the data flow.

(56) Many neural networks utilize computations such as element-wise calculation and depth-wise convolution for improved overall accuracy. Local processing elements are specialized for executing dense convolution with significant data reuse. The global buffer **208** may be utilized as the second-level data storage by local processing elements during dense convolution and may also perform computation for element-wise operations and depth-wise convolution. Global processing elements execute computation with low compute-to-memory ratio locally without communicating the data through layers (and hence, chips) of the neural network.

(57) The controller **608** may be local to each global processing element **602** or may be implemented by the chip master controller (controller **210**). Likewise, the global memory buffer **604** may be local to the global processing element **602** or implemented by the global buffer **208**.

(58) The neural network processor **300** leverages the locality of DNN algorithms to flexibly implement different data flows using a plurality of local processing elements and optionally one or

more global processing elements. Computation by the DNN layers may be tiled both spatially and temporally across multiple processing elements **206** to maximize utilization of the vector multiply-accumulate units **402**, maximize performance, minimize power consumed, or any balance of these factors, depending on the requirements of the implementation. The particular tiling and data flow may be determined by setting one or more of the configurable values discussed previously in conjunction with FIG. 5.

(59) FIG. 7 through FIG. 9 depict an embodiment of convolution **700** that may be utilized. This depiction is merely one example to illustrate the flexibility of the tiling possibilities.

(60) Generally, convolution of the input activations **702** with the weight tensors **704** is performed to generate the output activations **706**. The input activations **702** may have a number of input channels, C , as well as an input activation size, $H \times W$. The height, H , and the width, W , may be image pixel dimensions. The depth of the input to the multi-die package **202** (e.g., color depth of an image) may be the channel dimension of the input, C . Each of the weight tensors **704** may have C input channels and a weight kernel size, $R \times S$. The number of weight tensors **704** may equal the number of output channels, K . The output activations **706** may have a number of output channels, K , as well as output activation dimensions, $P \times Q$. In a multi-die package **202** embodiment, data may be distributed as follows: 1) weights are distributed between different dice along the input channel dimensions (C) and the output channel dimensions (K); 2) input activations are multicast along processing element rows of dice having a matching C dimension; and 3) output activations are reduced along processing element columns of dice. Within each die, weights are tiled similarly along the C and K dimensions of the processing elements **206**. Input activations are multicast via the network-on-a-chip router **212**. In addition to the C - K tiling illustrated in FIG. 7, the multi-die package **202** may also implement H - W tiling for network layers with a smaller number of input/output channels.

(61) Referring to FIG. 8, a chip-level tiling **800** comprises a chip input activation **802**, processing element input activations **804**, a chip **806**, processing element weights **808**, processing element output activations **810**, and a chip output activation **812**.

(62) The chip **806** may receive a number of processing element input activations **804**, from the chip input activation **802**. The number $C_{\text{sub.chip}}$ corresponds to the number of input channels the chip input activation **802**. The number $K_{\text{sub.chip}}$ corresponds to the number of output channels in the chip output activation **812**. Each column of PEs computes $K_{\text{sub.pe}}$ output channels to produce a tile of the output activation **810**. The chip input activation **802** is further divided along input channel dimensions into processing element input activations **804**, each with $C_{\text{sub.pe}}$ input channels, and applied to different rows of processing elements on the chip **806**, where they may be processed by $K_{\text{sub.chip}}/K_{\text{sub.pe}}$ processing elements to perform convolution operations in their vector multiply-accumulate units.

(63) The chip **806** has a number of processing elements (e.g., 16). Each processing element stores a portion of the processing element weights **808** with $C_{\text{sub.pe}}$ input channels and $K_{\text{sub.pe}}$ output channels, for example a portion corresponding to a particular convolution filter window.

(64) Each column of processing elements generates one of the processing element output activations **810**, each being a portion of the chip output activation **812**. The processing element output activations **810** are combined to generate the final chip output activation **812**. Each processing element may utilize a temporal processing element-level tiling **900** of the input activations as depicted in FIG. 9 to generate a portion of one of the processing element output activations **810**.

(65) Referring to FIG. 9, a processing element-level tiling **900** for a weight-stationary dataflow comprises a temporal tile **902**, a temporal tile **904**, a temporal tile **906**, a temporal tile **908**, a temporal tile **910**, and a temporal tile **912**—in other words, a temporal progression. Each tile of the temporal progression shows a pairing of input activations and weights for dot product/partial sum calculation by the vector MAC units in a particular PE at a different clock cycle. “Partial sum”

refers to an intermediate multiply-accumulate result in a dot-product-accumulate calculation.

(66) The multiply-accumulate units of the processing element compute a portion of a wide dot-product-accumulate as a partial result and store the partial result in the accumulation memory buffer. Partial results from multiply-accumulate units are stored and accumulated in the accumulation buffer, and once all the accumulations for an entry in the accumulation buffer complete, the accumulated value is forwarded to neighboring processing elements for cross-processing element accumulation (further reduction), or sent to the local post-processor for post-processing. “Accumulation memory buffer” refers to a memory buffer utilized to store computational results of one or more multiply-accumulate units. Partial results are transformed into a final result by the post-processor and communicated to the global processing element. The global processing element may stage the final results between layers of the deep neural network.

(67) The temporal tile **902** depicts a number of calculations performed at time, $T=0$. A first portion of the processing element input activation is convolved with a first portion of the processing element weights. The temporal tile **904** is a number of calculations performed at time, $T=1$. A second portion of the processing element input activation is convolved with the first portion of the processing element weights. This process continues for each portion of the processing element input activation until, in the temporal tile **906** at $T=n-1$, all portions of the processing element input activation have been convolved with a first portion of the processing element weights. The value “ n ” may correspond to the number of portions into which the processing element input activation is divided. In one embodiment, the processing element input activation may be divided into a number of rows and columns and each (row, column) iteration is calculated during one clock cycle, with a next iteration computed during the next clock cycle.

(68) Once the iterations for the processing element input activation is completed for the first portion of the processing element weights, the processing element-level tiling **900** iterates on the processing element input activation with a second portion of the processing element weights (the temporal tile **908**, the temporal tile **910**, and the temporal tile **912**). The temporal tile **908** is a number of calculations performed at time, $T=n$. The first portion of the processing element input activation is convolved with the second portion of the processing element weights. The temporal tile **910** is a number of calculations performed at time, $T=n+1$. The second portion of the processing element input activation is convolved with the second portion of the processing element weights. This process continues for each portion of the processing element input activation until, in the temporal tile **912** at $T=2n-1$, the last portion of the processing element input activation is convolved with the second portion of the processing element weights. The processing element input activation may be iterated similar to that as performed for the first portion of the processing element weights. The processing element-level tiling **900** may continue until each portion of processing element input activation is completed for each of the M portions of the processing element weights.

(69) FIG. **10** illustrates neural network processing **1000** on a multi-chip package in one embodiment. A global memory buffer allocation is made for each chip of the multi-chip package (block **1002**), and the initial weight values of the neural network are distributed to the weight buffers of each of the processing elements on each of the chip of the multi-chip package (block **1004**).

(70) An outer loop then executes for each neural network layer (opening loop block **1006**).

(71) The inner loop (opening loop block **1008**) executes the following for each tile: 1. load input activations from the global memory buffer to the level-one activation buffer of each processing element (block **1010**); 2. perform multiply-accumulate operations (block **1012**); 3. perform cross-processing element (intra-chip) accumulations (block **1014**); 4. perform cross-chip accumulations (block **1016**); 5. store results to global memory buffers (block **1018**).

(72) The global memory buffer allocation may be sequential or distributed. In a sequential allocation, input and output activations are allocated to consecutive memory addresses in the global

memory buffer. In a distributed allocation, input and output activations are interleaved in the global memory buffer to maximally exploit locality.

(73) FIG. 11 depicts an exemplary deeply-nested loop data flow **1100**. Table 1 below depicts loop reordering that may be implemented for different types of data flows.

(74) TABLE-US-00001 TABLE 1 L10 L11 L12 L20 L21 L22 L23 L24 L25 Loop Reordering WS 1 1 1 Q1 P1 C1 S R K1 L20-L21 can be reordered L22-L25 can be reordered OS 1 1 1 C1 S R Q1 P1 K1 L20-L22 can be reordered L23-L25 can be reordered IS 1 1 1 S R K1 C1 W1 H1 L20-L22 can be reordered L23-L25 can be reordered WS-LOS C0 R0 S0 Q1 P1 C1 S1 R1 K1 L20-L21 can be reordered L22-L25 can be reordered L10-L12 can be reordered OS-LWS Q0 P0 1 C1 S R Q1 P1 K1 L20-L22 can be reordered L23-L25 can be reordered L10-L11 can be reordered

(75) Additionally, an input activation collector can be used to improve input activation reuse along with weight and/or output reuse. For example when L20=R, L21=S, and L22=C1 . . . , input activations are accessed in a sliding window manner. This enables input activation reuse between consecutive iteration of L22 using an input activation collector (e.g., input collector **528**).

(76) FIG. 12-FIG. 16 depict various convolution computation data flows that may be implemented with embodiments of the described neural network acceleration architectures. A numeric postscript on a loop variable indicates the processing is carried out on the corresponding vector dimension by a particular processing element. The ordering of vector multiply-accumulate unit operations for the data flows is achieved by configuring the address generators to produce the correct sequence of addresses and applying them to the weight buffers and activation buffers of the processing elements. The amount of temporal reuse of operands of the multiply-accumulate units is controlled by the size of the various collectors. A larger collector size typically improves reuse in the collector registers, thereby reducing the energy consumption, but results in higher energy consumption in the collector registers themselves and can increase logic depth and timing sensitivity.

(77) FIG. 12 depicts a weight-stationary data flow **1200** (WS), in which the outer loops (lines 1-4) reuse the weights across dimensions (R, S, K1, C1) with different input activation vectors to execute vector multiply-accumulate operations for convolution calculation in the inner loop (lines 6-7). To implement a deep neural network accelerator with a WS data flow in FIG. 12, the depth of the weight collector (WC) may be set to 1 and the depth of the accumulation collector (AC) may set to 0 (zero). The amount of weight reuse is determined by dimensions (P1, Q1), which can be programmed (configured) at run-time. Weight stationary data flows provide configurable levels of reuse of weight values by the multiply-accumulate units.

(78) An output-stationary data flow **1300** (OS) is depicted in FIG. 13. An OS data flow for convolution computation may be implemented on the deep neural network accelerator by configuring the depth of WC to 0 (zero) and the depth of the AC to 1 (one). The amount of output reuse is determined by dimensions (R, S, C1), which can be programmed at run-time. Output stationary data flows provide configurable levels of reuse of computational results of the multiply-accumulate units.

(79) An input-stationary data flow **1400** (IS) is depicted in FIG. 14. An IS data flow may be implemented for convolution computation by the deep neural network accelerator by configuring the depth of WC to 0 (zero) and the depth of the AC also to 0 (zero). The depth of the input collector (IC) is set to 1 (one). The amount of input reuse is determined by dimensions (K1, R, S), which can be programmed at run-time. Input stationary data flows provide configurable levels of reuse of input activation values by the multiply-accumulate units during convolution computation. The input-stationary data flow **1400** may also be referred to as an activation stationary data flow or convolution.

(80) FIG. 15 depicts a weight-stationary-local-output-stationary data flow **1500** (WS-LOS) data flow, a type of multi-level WS-OS data flow for convolution computation. Unlike conventional data flows that are optimized for reuse of one operand only, multi-level data flows of this type can reuse both weights and outputs to configurable degrees. For example the weight-stationary-local-

output-stationary data flow **1500** may reuse multiply-accumulate unit outputs in the innermost loop in a single-entry accumulation collector ($AC=1$), while reusing weights in the outer loops via a multi-entry weight collector ($WC=C0>1$). As depicted in the weight-stationary-local-output-stationary data flow **1500**, outputs of the vector multiply-accumulate units are reused in the $C0$ loop while weights are reused in the $P1$ and $Q1$ loops.

(81) FIG. **16** depicts an output-stationary-local-weight-stationary data flow **1600** (OS-LWS), another types of multi-level WS-OS convolution data flow. The OS-LWS data flow prioritizes weight reuse in the innermost loop (line 9) while providing vector multiply-accumulate unit output reuse in the outer loops (lines 5-7). In the depicted example, a single-entry weight collector and a multiple-entry accumulation collector of depth ($AC=Q0>1$) are configured to implement the output-stationary-local-weight-stationary data flow **1600**.

(82) Other multi-level data flow examples will be readily apparent to those of ordinary skill in the art based on these examples.

(83) A neural network accelerator may thus be configured with logic (e.g., a memory configured with instructions, registers configured with settings) to configure a depth of the weight collection buffer, a depth of the accumulation collection buffer, and/or a depth of an input activation collection buffer to implement a deep neural network accelerator with any of a weight stationary convolution of the weights and activations, an output stationary convolution of the weights and activations, an activation stationary convolution of the weights and activations, and combinations of these data flow types.

(84) The mechanisms disclosed herein may be implemented by computing devices utilizing one or more graphic processing unit (GPU) and/or general-purpose data processor (e.g., a 'central processing unit or CPU). Exemplary architectures will now be described that may be configured to carry out the techniques disclosed herein on such devices.

(85) Referring to FIG. **17**-FIG. **22**, in one example, the input activations and weights are stored in a register file, and/or shared/global memory such as shared memory/L1 cache **2018**, level two cache **1904**, and/or memory **1704**. This takes the place of the input activation buffer and weight buffer of the processing elements. Partial sum results may also be stored in a register file, and/or the shared memory/L1 cache **2018**, level two cache **1904**, and/or memory **1704**. The vector multiply-accumulate units may be implemented in the tensor cores (see the description in conjunction with FIG. **20**). The global buffer may be implemented for example by the level two cache **1904**, or memory **1704**.

(86) One or more tensor cores are included in the core **2010** modules. The various collectors (e.g., weight collector, accumulation collector, and input collector) may be implemented in the tensor cores where the vector multiply-accumulate units are disposed to implement various dataflows. As explained before, WS, OS, and IS dataflows may utilize only a single-entry of weight collector, accumulation collector, and input collector respectively. Multi-level dataflow execution utilizes multiple collectors for data staging and reuse. For example, an OS-LWS dataflow utilizes a single-entry weight collector and multiple-entry accumulation collector. In addition, input collectors may be used to capture reuse of input activations. Both convolution and MMA operations may utilize multi-level dataflows to improve energy and/or execution efficiency. The register file **2008** (or files) or shared memory/L1 cache may operate as a weight buffer, input buffer, or accumulation buffer. One or more streaming multiprocessor **2000** may operate as the local processing element controller **410** and/or global processing element controller **608** to provide data orchestration across different levels of memory, and to provide sequencing of operations. The various address generators and buffer managers may be implemented for example by the M-pipe controller **1810** and/or memory management unit **1816** and may be configured based with instructions to generate a sequence of read/write addresses for input activations, weights and partial sums. The overall data paths may be carried out using parallel processing unit (see parallel processing unit **1700**) operations performed by one or more streaming multiprocessor (see streaming multiprocessor

2000) data path.

(87) The following description may use certain acronyms and abbreviations as follows: “DPC” refers to a “data processing cluster”; “GPC” refers to a “general processing cluster”; “I/O” refers to a “input/output”; “L1 cache” refers to “level one cache”; “L2 cache” refers to “level two cache”; “LSU” refers to a “load/store unit”; “MMU” refers to a “memory management unit”; “MPC” refers to an “M-pipe controller”; “PPU” refers to a “parallel processing unit”; “PROP” refers to a “pre-raster operations unit”; “ROP” refers to a “raster operations”; “SFU” refers to a “special function unit”; “SM” refers to a “streaming multiprocessor”; “Viewport SCC” refers to “viewport scale, cull, and clip”; “WDX” refers to a “work distribution crossbar”; and. “XBar” refers to a “crossbar”.

Parallel Processing Unit

(88) FIG. 17 depicts a parallel processing unit **1700**, in accordance with an embodiment. In an embodiment, the parallel processing unit **1700** is a multi-threaded processor that is implemented on one or more integrated circuit devices. The parallel processing unit **1700** is a latency hiding architecture designed to process many threads in parallel. A thread (e.g., a thread of execution) is an instantiation of a set of instructions configured to be executed by the parallel processing unit **1700**. In an embodiment, the parallel processing unit **1700** is a graphics processing unit (GPU) configured to implement a graphics rendering pipeline for processing three-dimensional (3D) graphics data in order to generate two-dimensional (2D) image data for display on a display device such as a liquid crystal display (LCD) device. In other embodiments, the parallel processing unit **1700** may be utilized for performing general-purpose computations. While one exemplary parallel processor is provided herein for illustrative purposes, it should be strongly noted that such processor is set forth for illustrative purposes only, and that any processor may be employed to supplement and/or substitute for the same.

(89) One or more parallel processing unit **1700** modules may be configured to accelerate thousands of High Performance Computing (HPC), data center, and machine learning applications. The parallel processing unit **1700** may be configured to accelerate numerous deep learning systems and applications including autonomous vehicle platforms, deep learning, high-accuracy speech, image, and text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and the like.

(90) As shown in FIG. 17, the parallel processing unit **1700** includes an I/O unit **1706**, a front-end unit **1710**, a scheduler unit **1712**, a work distribution unit **1714**, a hub **1716**, a crossbar **1718**, one or more general processing cluster **1800** modules, and one or more memory partition unit **1900** modules. The parallel processing unit **1700** may be connected to a host processor or other parallel processing unit **1700** modules via one or more high-speed NVLink **1708** interconnects. The parallel processing unit **1700** may be connected to a host processor or other peripheral devices via an interconnect **1702**. The parallel processing unit **1700** may also be connected to a local memory comprising a number of memory **1704** devices. In an embodiment, the local memory may comprise a number of dynamic random access memory (DRAM) devices. The DRAM devices may be configured as a high-bandwidth memory (HBM) subsystem, with multiple DRAM dies stacked within each device. The memory **1704** may comprise logic to configure the parallel processing unit **1700** to carry out aspects of the techniques disclosed herein.

(91) The NVLink **1708** interconnect enables systems to scale and include one or more parallel processing unit **1700** modules combined with one or more CPUs, supports cache coherence between the parallel processing unit **1700** modules and CPUs, and CPU mastering. Data and/or commands may be transmitted by the NVLink **1708** through the hub **1716** to/from other units of the parallel processing unit **1700** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). The NVLink **1708** is described in more detail in conjunction with FIG. 21.

(92) The I/O unit **1706** is configured to transmit and receive communications (e.g., commands, data, etc.) from a host processor (not shown) over the interconnect **1702**. The I/O unit **1706** may communicate with the host processor directly via the interconnect **1702** or through one or more intermediate devices such as a memory bridge. In an embodiment, the I/O unit **1706** may communicate with one or more other processors, such as one or more parallel processing unit **1700** modules via the interconnect **1702**. In an embodiment, the I/O unit **1706** implements a Peripheral Component Interconnect Express (PCIe) interface for communications over a PCIe bus and the interconnect **1702** is a PCIe bus. In alternative embodiments, the I/O unit **1706** may implement other types of well-known interfaces for communicating with external devices.

(93) The I/O unit **1706** decodes packets received via the interconnect **1702**. In an embodiment, the packets represent commands configured to cause the parallel processing unit **1700** to perform various operations. The I/O unit **1706** transmits the decoded commands to various other units of the parallel processing unit **1700** as the commands may specify. For example, some commands may be transmitted to the front-end unit **1710**. Other commands may be transmitted to the hub **1716** or other units of the parallel processing unit **1700** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). In other words, the I/O unit **1706** is configured to route communications between and among the various logical units of the parallel processing unit **1700**.

(94) In an embodiment, a program executed by the host processor encodes a command stream in a buffer that provides workloads to the parallel processing unit **1700** for processing. A workload may comprise several instructions and data to be processed by those instructions. The buffer is a region in a memory that is accessible (e.g., read/write) by both the host processor and the parallel processing unit **1700**. For example, the I/O unit **1706** may be configured to access the buffer in a system memory connected to the interconnect **1702** via memory requests transmitted over the interconnect **1702**. In an embodiment, the host processor writes the command stream to the buffer and then transmits a pointer to the start of the command stream to the parallel processing unit **1700**. The front-end unit **1710** receives pointers to one or more command streams. The front-end unit **1710** manages the one or more streams, reading commands from the streams and forwarding commands to the various units of the parallel processing unit **1700**.

(95) The front-end unit **1710** is coupled to a scheduler unit **1712** that configures the various general processing cluster **1800** modules to process tasks defined by the one or more streams. The scheduler unit **1712** is configured to track state information related to the various tasks managed by the scheduler unit **1712**. The state may indicate which general processing cluster **1800** a task is assigned to, whether the task is active or inactive, a priority level associated with the task, and so forth. The scheduler unit **1712** manages the execution of a plurality of tasks on the one or more general processing cluster **1800** modules.

(96) The scheduler unit **1712** is coupled to a work distribution unit **1714** that is configured to dispatch tasks for execution on the general processing cluster **1800** modules. The work distribution unit **1714** may track a number of scheduled tasks received from the scheduler unit **1712**. In an embodiment, the work distribution unit **1714** manages a pending task pool and an active task pool for each of the general processing cluster **1800** modules. The pending task pool may comprise a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular general processing cluster **1800**. The active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by the general processing cluster **1800** modules. As a general processing cluster **1800** finishes the execution of a task, that task is evicted from the active task pool for the general processing cluster **1800** and one of the other tasks from the pending task pool is selected and scheduled for execution on the general processing cluster **1800**. If an active task has been idle on the general processing cluster **1800**, such as while waiting for a data dependency to be resolved, then the active task may be evicted from the general processing cluster **1800** and returned to the pending task pool while another task in the pending task pool is selected

and scheduled for execution on the general processing cluster **1800**.

(97) The work distribution unit **1714** communicates with the one or more general processing cluster **1800** modules via crossbar **1718**. The crossbar **1718** is an interconnect network that couples many of the units of the parallel processing unit **1700** to other units of the parallel processing unit **1700**. For example, the crossbar **1718** may be configured to couple the work distribution unit **1714** to a particular general processing cluster **1800**. Although not shown explicitly, one or more other units of the parallel processing unit **1700** may also be connected to the crossbar **1718** via the hub **1716**.

(98) The tasks are managed by the scheduler unit **1712** and dispatched to a general processing cluster **1800** by the work distribution unit **1714**. The general processing cluster **1800** is configured to process the task and generate results. The results may be consumed by other tasks within the general processing cluster **1800**, routed to a different general processing cluster **1800** via the crossbar **1718**, or stored in the memory **1704**. The results can be written to the memory **1704** via the memory partition unit **1900** modules, which implement a memory interface for reading and writing data to/from the memory **1704**. The results can be transmitted to another parallel processing unit **1700** or CPU via the NVLink **1708**. In an embodiment, the parallel processing unit **1700** includes a number U of memory partition unit **1900** modules that is equal to the number of separate and distinct memory **1704** devices coupled to the parallel processing unit **1700**. A memory partition unit **1900** will be described in more detail below in conjunction with FIG. **19**.

(99) In an embodiment, a host processor executes a driver kernel that implements an application programming interface (API) that enables one or more applications executing on the host processor to schedule operations for execution on the parallel processing unit **1700**. In an embodiment, multiple compute applications are simultaneously executed by the parallel processing unit **1700** and the parallel processing unit **1700** provides isolation, quality of service (QoS), and independent address spaces for the multiple compute applications. An application may generate instructions (e.g., API calls) that cause the driver kernel to generate one or more tasks for execution by the parallel processing unit **1700**. The driver kernel outputs tasks to one or more streams being processed by the parallel processing unit **1700**. Each task may comprise one or more groups of related threads, referred to herein as a warp. In an embodiment, a warp comprises 32 related threads that may be executed in parallel. Cooperating threads may refer to a plurality of threads including instructions to perform the task and that may exchange data through shared memory. Threads and cooperating threads are described in more detail in conjunction with FIG. **20**.

(100) FIG. **18** depicts a general processing cluster **1800** of the parallel processing unit **1700** of FIG. **17**, in accordance with an embodiment. As shown in FIG. **18**, each general processing cluster **1800** includes a number of hardware units for processing tasks. In an embodiment, each general processing cluster **1800** includes a pipeline manager **1802**, a pre-raster operations unit **1804**, a raster engine **1808**, a work distribution crossbar **1814**, a memory management unit **1816**, and one or more data processing cluster **1806**. It will be appreciated that the general processing cluster **1800** of FIG. **18** may include other hardware units in lieu of or in addition to the units shown in FIG. **18**.

(101) In an embodiment, the operation of the general processing cluster **1800** is controlled by the pipeline manager **1802**. The pipeline manager **1802** manages the configuration of the one or more data processing cluster **1806** modules for processing tasks allocated to the general processing cluster **1800**. In an embodiment, the pipeline manager **1802** may configure at least one of the one or more data processing cluster **1806** modules to implement at least a portion of a graphics rendering pipeline. For example, a data processing cluster **1806** may be configured to execute a vertex shader program on the programmable streaming multiprocessor **2000**. The pipeline manager **1802** may also be configured to route packets received from the work distribution unit **1714** to the appropriate logical units within the general processing cluster **1800**. For example, some packets may be routed to fixed function hardware units in the pre-raster operations unit **1804** and/or raster engine **1808** while other packets may be routed to the data processing cluster **1806** modules for processing by the primitive engine **1812** or the streaming multiprocessor **2000**. In an embodiment,

the pipeline manager **1802** may configure at least one of the one or more data processing cluster **1806** modules to implement a neural network model and/or a computing pipeline.

(102) The pre-raster operations unit **1804** is configured to route data generated by the raster engine **1808** and the data processing cluster **1806** modules to a Raster Operations (ROP) unit, described in more detail in conjunction with FIG. **19**. The pre-raster operations unit **1804** may also be configured to perform optimizations for color blending, organize pixel data, perform address translations, and the like.

(103) The raster engine **1808** includes a number of fixed function hardware units configured to perform various raster operations. In an embodiment, the raster engine **1808** includes a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, and a tile coalescing engine. The setup engine receives transformed vertices and generates plane equations associated with the geometric primitive defined by the vertices. The plane equations are transmitted to the coarse raster engine to generate coverage information (e.g., an x, y coverage mask for a tile) for the primitive. The output of the coarse raster engine is transmitted to the culling engine where fragments associated with the primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. Those fragments that survive clipping and culling may be passed to the fine raster engine to generate attributes for the pixel fragments based on the plane equations generated by the setup engine. The output of the raster engine **1808** comprises fragments to be processed, for example, by a fragment shader implemented within a data processing cluster **1806**.

(104) Each data processing cluster **1806** included in the general processing cluster **1800** includes an M-pipe controller **1810**, a primitive engine **1812**, and one or more streaming multiprocessor **2000** modules. The M-pipe controller **1810** controls the operation of the data processing cluster **1806**, routing packets received from the pipeline manager **1802** to the appropriate units in the data processing cluster **1806**. For example, packets associated with a vertex may be routed to the primitive engine **1812**, which is configured to fetch vertex attributes associated with the vertex from the memory **1704**. In contrast, packets associated with a shader program may be transmitted to the streaming multiprocessor **2000**.

(105) The streaming multiprocessor **2000** comprises a programmable streaming processor that is configured to process tasks represented by a number of threads. Each streaming multiprocessor **2000** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently. In an embodiment, the streaming multiprocessor **2000** implements a Single-Instruction, Multiple-Data (SIMD) architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on the same set of instructions. All threads in the group of threads execute the same instructions. In another embodiment, the streaming multiprocessor **2000** implements a Single-Instruction, Multiple Thread (SIMT) architecture where each thread in a group of threads is configured to process a different set of data based on the same set of instructions, but where individual threads in the group of threads are allowed to diverge during execution. In an embodiment, a program counter, call stack, and execution state is maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within the warp diverge. In another embodiment, a program counter, call stack, and execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. When execution state is maintained for each individual thread, threads executing the same instructions may be converged and executed in parallel for maximum efficiency. The streaming multiprocessor **2000** will be described in more detail below in conjunction with FIG. **20**.

(106) The memory management unit **1816** provides an interface between the general processing cluster **1800** and the memory partition unit **1900**. The memory management unit **1816** may provide translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In an embodiment, the memory management unit **1816** provides one or more

translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in the memory **1704**.

(107) FIG. **19** depicts a memory partition unit **1900** of the parallel processing unit **1700** of FIG. **17**, in accordance with an embodiment. As shown in FIG. **19**, the memory partition unit **1900** includes a raster operations unit **1902**, a level two cache **1904**, and a memory interface **1906**. The memory interface **1906** is coupled to the memory **1704**. Memory interface **1906** may implement 32, 64, 128, 1024-bit data buses, or the like, for high-speed data transfer. In an embodiment, the parallel processing unit **1700** incorporates U memory interface **1906** modules, one memory interface **1906** per pair of memory partition unit **1900** modules, where each pair of memory partition unit **1900** modules is connected to a corresponding memory **1704** device. For example, parallel processing unit **1700** may be connected to up to Y memory **1704** devices, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random access memory, or other types of persistent storage.

(108) In an embodiment, the memory interface **1906** implements an HBM2 memory interface and Y equals half U. In an embodiment, the HBM2 memory stacks are located on the same physical package as the parallel processing unit **1700**, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In an embodiment, each HBM2 stack includes four memory dies and Y equals 4, with HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits.

(109) In an embodiment, the memory **1704** supports Single-Error Correcting Double-Error Detecting (SECDED) Error Correction Code (ECC) to protect data. ECC provides higher reliability for compute applications that are sensitive to data corruption. Reliability is especially important in large-scale cluster computing environments where parallel processing unit **1700** modules process very large datasets and/or run applications for extended periods.

(110) In an embodiment, the parallel processing unit **1700** implements a multi-level memory hierarchy. In an embodiment, the memory partition unit **1900** supports a unified memory to provide a single unified virtual address space for CPU and parallel processing unit **1700** memory, enabling data sharing between virtual memory systems. In an embodiment the frequency of accesses by a parallel processing unit **1700** to memory located on other processors is traced to ensure that memory pages are moved to the physical memory of the parallel processing unit **1700** that is accessing the pages more frequently. In an embodiment, the NVLink **1708** supports address translation services allowing the parallel processing unit **1700** to directly access a CPU's page tables and providing full access to CPU memory by the parallel processing unit **1700**.

(111) In an embodiment, copy engines transfer data between multiple parallel processing unit **1700** modules or between parallel processing unit **1700** modules and CPUs. The copy engines can generate page faults for addresses that are not mapped into the page tables. The memory partition unit **1900** can then service the page faults, mapping the addresses into the page table, after which the copy engine can perform the transfer. In a conventional system, memory is pinned (e.g., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing the available memory. With hardware page faulting, addresses can be passed to the copy engines without worrying if the memory pages are resident, and the copy process is transparent.

(112) Data from the memory **1704** or other system memory may be fetched by the memory partition unit **1900** and stored in the level two cache **1904**, which is located on-chip and is shared between the various general processing cluster **1800** modules. As shown, each memory partition unit **1900** includes a portion of the level two cache **1904** associated with a corresponding memory **1704** device. Lower level caches may then be implemented in various units within the general processing cluster **1800** modules. For example, each of the streaming multiprocessor **2000** modules may implement an L1 cache. The L1 cache is private memory that is dedicated to a particular streaming multiprocessor **2000**. Data from the level two cache **1904** may be fetched and stored in each of the L1 caches for processing in the functional units of the streaming multiprocessor **2000**.

modules. The level two cache **1904** is coupled to the memory interface **1906** and the crossbar **1718**. (113) The raster operations unit **1902** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and the like. The raster operations unit **1902** also implements depth testing in conjunction with the raster engine **1808**, receiving a depth for a sample location associated with a pixel fragment from the culling engine of the raster engine **1808**. The depth is tested against a corresponding depth in a depth buffer for a sample location associated with the fragment. If the fragment passes the depth test for the sample location, then the raster operations unit **1902** updates the depth buffer and transmits a result of the depth test to the raster engine **1808**. It will be appreciated that the number of partition memory partition unit **1900** modules may be different than the number of general processing cluster **1800** modules and, therefore, each raster operations unit **1902** may be coupled to each of the general processing cluster **1800** modules. The raster operations unit **1902** tracks packets received from the different general processing cluster **1800** modules and determines which general processing cluster **1800** that a result generated by the raster operations unit **1902** is routed to through the crossbar **1718**. Although the raster operations unit **1902** is included within the memory partition unit **1900** in FIG. **19**, in other embodiment, the raster operations unit **1902** may be outside of the memory partition unit **1900**. For example, the raster operations unit **1902** may reside in the general processing cluster **1800** or another unit.

(114) FIG. **20** illustrates the streaming multiprocessor **2000** of FIG. **18**, in accordance with an embodiment. As shown in FIG. **20**, the streaming multiprocessor **2000** includes an instruction cache **2002**, one or more scheduler unit **2004** modules (e.g., such as scheduler unit **1712**), a register file **2008**, one or more processing core **2010** modules, one or more special function unit **2012** modules, one or more load/store unit **2014** modules, an interconnect network **2016**, and a shared memory/L1 cache **2018**.

(115) As described above, the work distribution unit **1714** dispatches tasks for execution on the general processing cluster **1800** modules of the parallel processing unit **1700**. The tasks are allocated to a particular data processing cluster **1806** within a general processing cluster **1800** and, if the task is associated with a shader program, the task may be allocated to a streaming multiprocessor **2000**. The scheduler unit **1712** receives the tasks from the work distribution unit **1714** and manages instruction scheduling for one or more thread blocks assigned to the streaming multiprocessor **2000**. The scheduler unit **2004** schedules thread blocks for execution as warps of parallel threads, where each thread block is allocated at least one warp. In an embodiment, each warp executes 32 threads. The scheduler unit **2004** may manage a plurality of different thread blocks, allocating the warps to the different thread blocks and then dispatching instructions from the plurality of different cooperative groups to the various functional units (e.g., core **2010** modules, special function unit **2012** modules, and load/store unit **2014** modules) during each clock cycle.

(116) Cooperative Groups is a programming model for organizing groups of communicating threads that allows developers to express the granularity at which threads are communicating, enabling the expression of richer, more efficient parallel decompositions. Cooperative launch APIs support synchronization amongst thread blocks for the execution of parallel algorithms. Conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., the `syncthreads ( )` function). However, programmers would often like to define groups of threads at smaller than thread block granularities and synchronize within the defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces.

(117) Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (e.g., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on the threads in a cooperative group. The programming model supports clean composition across software boundaries, so that libraries and utility functions can

synchronize safely within their local context without having to make assumptions about convergence. Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

(118) A dispatch **2006** unit is configured within the scheduler unit **2004** to transmit instructions to one or more of the functional units. In one embodiment, the scheduler unit **2004** includes two dispatch **2006** units that enable two different instructions from the same warp to be dispatched during each clock cycle. In alternative embodiments, each scheduler unit **2004** may include a single dispatch **2006** unit or additional dispatch **2006** units.

(119) Each streaming multiprocessor **2000** includes a register file **2008** that provides a set of registers for the functional units of the streaming multiprocessor **2000**. In an embodiment, the register file **2008** is divided between each of the functional units such that each functional unit is allocated a dedicated portion of the register file **2008**. In another embodiment, the register file **2008** is divided between the different warps being executed by the streaming multiprocessor **2000**. The register file **2008** provides temporary storage for operands connected to the data paths of the functional units.

(120) Each streaming multiprocessor **2000** comprises L processing core **2010** modules. In an embodiment, the streaming multiprocessor **2000** includes a large number (e.g., 128, etc.) of distinct processing core **2010** modules. Each core **2010** may include a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes a floating point arithmetic logic unit and an integer arithmetic logic unit. In an embodiment, the floating point arithmetic logic units implement the IEEE 754-2008 standard for floating point arithmetic. In an embodiment, the core **2010** modules include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

(121) Tensor cores configured to perform matrix operations, and, in an embodiment, one or more tensor cores are included in the core **2010** modules. In particular, the tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In an embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation $D = A'B + C$, where A, B, C, and D are 4×4 matrices.

(122) In an embodiment, the matrix multiply inputs A and B are 16-bit floating point matrices, while the accumulation matrices C and D may be 16-bit floating point or 32-bit floating point matrices. Tensor Cores operate on 16-bit floating point input data with 32-bit floating point accumulation. The 16-bit floating point multiply requires 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with the other intermediate products for a $4 \times 4 \times 4$ matrix multiply. In practice, Tensor Cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements. An API, such as CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use Tensor Cores from a CUDA-C++ program. At the CUDA level, the warp-level interface assumes 16×16 size matrices spanning all 32 threads of the warp.

(123) Each streaming multiprocessor **2000** also comprises M special function unit **2012** modules that perform special functions (e.g., attribute evaluation, reciprocal square root, and the like). In an embodiment, the special function unit **2012** modules may include a tree traversal unit configured to traverse a hierarchical tree data structure. In an embodiment, the special function unit **2012** modules may include texture unit configured to perform texture map filtering operations. In an embodiment, the texture units are configured to load texture maps (e.g., a 2D array of texels) from the memory **1704** and sample the texture maps to produce sampled texture values for use in shader programs executed by the streaming multiprocessor **2000**. In an embodiment, the texture maps are stored in the shared memory/L1 cache **2018**. The texture units implement texture operations such

as filtering operations using mip-maps (e.g., texture maps of varying levels of detail). In an embodiment, each streaming multiprocessor **2000** includes two texture units.

(124) Each streaming multiprocessor **2000** also comprises N load/store unit **2014** modules that implement load and store operations between the shared memory/L1 cache **2018** and the register file **2008**. Each streaming multiprocessor **2000** includes an interconnect network **2016** that connects each of the functional units to the register file **2008** and the load/store unit **2014** to the register file **2008** and shared memory/L1 cache **2018**. In an embodiment, the interconnect network **2016** is a crossbar that can be configured to connect any of the functional units to any of the registers in the register file **2008** and connect the load/store unit **2014** modules to the register file **2008** and memory locations in shared memory/L1 cache **2018**.

(125) The shared memory/L1 cache **2018** is an array of on-chip memory that allows for data storage and communication between the streaming multiprocessor **2000** and the primitive engine **1812** and between threads in the streaming multiprocessor **2000**. In an embodiment, the shared memory/L1 cache **2018** comprises 128 KB of storage capacity and is in the path from the streaming multiprocessor **2000** to the memory partition unit **1900**. The shared memory/L1 cache **2018** can be used to cache reads and writes. One or more of the shared memory/L1 cache **2018**, level two cache **1904**, and memory **1704** are backing stores.

(126) Combining data cache and shared memory functionality into a single memory block provides the best overall performance for both types of memory accesses. The capacity is usable as a cache by programs that do not use shared memory. For example, if shared memory is configured to use half of the capacity, texture and load/store operations can use the remaining capacity. Integration within the shared memory/L1 cache **2018** enables the shared memory/L1 cache **2018** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data.

(127) When configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. Specifically, the fixed function graphics processing units shown in FIG. **17**, are bypassed, creating a much simpler programming model. In the general purpose parallel computation configuration, the work distribution unit **1714** assigns and distributes blocks of threads directly to the data processing cluster **1806** modules. The threads in a block execute the same program, using a unique thread ID in the calculation to ensure each thread generates unique results, using the streaming multiprocessor **2000** to execute the program and perform calculations, shared memory/L1 cache **2018** to communicate between threads, and the load/store unit **2014** to read and write global memory through the shared memory/L1 cache **2018** and the memory partition unit **1900**. When configured for general purpose parallel computation, the streaming multiprocessor **2000** can also write commands that the scheduler unit **1712** can use to launch new work on the data processing cluster **1806** modules.

(128) The parallel processing unit **1700** may be included in a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and the like. In an embodiment, the parallel processing unit **1700** is embodied on a single semiconductor substrate. In another embodiment, the parallel processing unit **1700** is included in a system-on-a-chip (SoC) along with one or more other devices such as additional parallel processing unit **1700** modules, the memory **1704**, a reduced instruction set computer (RISC) CPU, a memory management unit (MMU), a digital-to-analog converter (DAC), and the like.

(129) In an embodiment, the parallel processing unit **1700** may be included on a graphics card that includes one or more memory devices **304**. The graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In yet another embodiment, the parallel processing unit **1700** may be an integrated graphics processing unit (iGPU) or parallel processor included in the chipset of the motherboard.

(130) Exemplary Computing System

(131) Systems with multiple GPUs and CPUs are used in a variety of industries as developers expose and leverage more parallelism in applications such as artificial intelligence computing. High-performance GPU-accelerated systems with tens to many thousands of compute nodes are deployed in data centers, research facilities, and supercomputers to solve ever larger problems. As the number of processing devices within the high-performance systems increases, the communication and data transfer mechanisms need to scale to support the increased bandwidth.

(132) FIG. 21 is a conceptual diagram of a processing system 2100 implemented using the parallel processing unit 1700 of FIG. 17, in accordance with an embodiment. The processing system 2100 includes a central processing unit 2106, switch 2102, and multiple parallel processing unit 1700 modules each and respective memory 1704 modules. The NVLink 1708 provides high-speed communication links between each of the parallel processing unit 1700 modules. Although a particular number of NVLink 1708 and interconnect 1702 connections are illustrated in FIG. 21, the number of connections to each parallel processing unit 1700 and the central processing unit 2106 may vary. The switch 2102 interfaces between the interconnect 1702 and the central processing unit 2106. The parallel processing unit 1700 modules, memory 1704 modules, and NVLink 1708 connections may be situated on a single semiconductor platform to form a parallel processing module 2104. In an embodiment, the switch 2102 supports two or more protocols to interface between various different connections and/or links.

(133) In another embodiment (not shown), the NVLink 1708 provides one or more high-speed communication links between each of the parallel processing unit 1700 modules and the central processing unit 2106 and the switch 2102 interfaces between the interconnect 1702 and each of the parallel processing unit 1700 modules. The parallel processing unit 1700 modules, memory 1704 modules, and interconnect 1702 may be situated on a single semiconductor platform to form a parallel processing module 2104. In yet another embodiment (not shown), the interconnect 1702 provides one or more communication links between each of the parallel processing unit 1700 modules and the central processing unit 2106 and the switch 2102 interfaces between each of the parallel processing unit 1700 modules using the NVLink 1708 to provide one or more high-speed communication links between the parallel processing unit 1700 modules. In another embodiment (not shown), the NVLink 1708 provides one or more high-speed communication links between the parallel processing unit 1700 modules and the central processing unit 2106 through the switch 2102. In yet another embodiment (not shown), the interconnect 1702 provides one or more communication links between each of the parallel processing unit 1700 modules directly. One or more of the NVLink 1708 high-speed communication links may be implemented as a physical NVLink interconnect or either an on-chip or on-die interconnect using the same protocol as the NVLink 1708.

(134) In the context of the present description, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit fabricated on a die or chip. It should be noted that the term single semiconductor platform may also refer to multi-chip modules with increased connectivity which simulate on-chip operation and make substantial improvements over utilizing a conventional bus implementation. Of course, the various circuits or devices may also be situated separately or in various combinations of semiconductor platforms per the desires of the user. Alternately, the parallel processing module 2104 may be implemented as a circuit board substrate and each of the parallel processing unit 1700 modules and/or memory 1704 modules may be packaged devices. In an embodiment, the central processing unit 2106, switch 2102, and the parallel processing module 2104 are situated on a single semiconductor platform.

(135) In an embodiment, the signaling rate of each NVLink 1708 is 20 to 25 Gigabits/second and each parallel processing unit 1700 includes six NVLink 1708 interfaces (as shown in FIG. 21, five NVLink 1708 interfaces are included for each parallel processing unit 1700). Each NVLink 1708 provides a data transfer rate of 25 Gigabytes/second in each direction, with six links providing 300

Gigabytes/second. The NVLink **1708** can be used exclusively for PPU-to-PPU communication as shown in FIG. **21**, or some combination of PPU-to-PPU and PPU-to-CPU, when the central processing unit **2106** also includes one or more NVLink **1708** interfaces.

(136) In an embodiment, the NVLink **1708** allows direct load/store/atomic access from the central processing unit **2106** to each parallel processing unit **1700** module's memory **1704**. In an embodiment, the NVLink **1708** supports coherency operations, allowing data read from the memory **1704** modules to be stored in the cache hierarchy of the central processing unit **2106**, reducing cache access latency for the central processing unit **2106**. In an embodiment, the NVLink **1708** includes support for Address Translation Services (ATS), allowing the parallel processing unit **1700** to directly access page tables within the central processing unit **2106**. One or more of the NVLink **1708** may also be configured to operate in a low-power mode.

(137) FIG. **22** depicts an exemplary processing system **2200** in which the various architecture and/or functionality of the various previous embodiments may be implemented. As shown, an exemplary processing system **2200** is provided including at least one central processing unit **2106** that is connected to a communications bus **2210**. The communication communications bus **2210** may be implemented using any suitable protocol, such as PCI (Peripheral Component Interconnect), PCI-Express, AGP (Accelerated Graphics Port), HyperTransport, or any other bus or point-to-point communication protocol(s). The exemplary processing system **2200** also includes a main memory **2204**. Control logic (software) and data are stored in the main memory **2204** which may take the form of random access memory (RAM).

(138) The exemplary processing system **2200** also includes input devices **2208**, the parallel processing module **2104**, and display devices **2206**, e.g. a conventional CRT (cathode ray tube), LCD (liquid crystal display), LED (light emitting diode), plasma display or the like. User input may be received from the input devices **2208**, e.g., keyboard, mouse, touchpad, microphone, and the like. Each of the foregoing modules and/or devices may even be situated on a single semiconductor platform to form the exemplary processing system **2200**. Alternately, the various modules may also be situated separately or in various combinations of semiconductor platforms per the desires of the user.

(139) Further, the exemplary processing system **2200** may be coupled to a network (e.g., a telecommunications network, local area network (LAN), wireless network, wide area network (WAN) such as the Internet, peer-to-peer network, cable network, or the like) through a network interface **2202** for communication purposes.

(140) The exemplary processing system **2200** may also include a secondary storage (not shown). The secondary storage includes, for example, a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (DVD) drive, recording device, universal serial bus (USB) flash memory. The removable storage drive reads from and/or writes to a removable storage unit in a well-known manner.

(141) Computer programs, or computer control logic algorithms, may be stored in the main memory **2204** and/or the secondary storage. Such computer programs, when executed, enable the exemplary processing system **2200** to perform various functions. The main memory **2204**, the storage, and/or any other storage are possible examples of computer-readable media.

(142) The architecture and/or functionality of the various previous figures may be implemented in the context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and/or any other desired system. For example, the exemplary processing system **2200** may take the form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

(143) While various embodiments have been described above, it should be understood that they

have been presented by way of example only, and not limitation. Thus, the breadth and scope of a preferred embodiment should not be limited by any of the above-described exemplary embodiments, but should be defined only in accordance with the following claims and their equivalents.

(144) Various logic functional operations described herein may be implemented in logic that is referred to using a noun or noun phrase reflecting said operation or function. For example, an association operation may be carried out by an “associator” or “correlator”. Likewise, switching may be carried out by a “switch”, selection by a “selector”, and so on.

(145) “Logic” is used herein to machine memory circuits, non transitory machine readable media, and/or circuitry which by way of its material and/or material-energy configuration comprises control and/or procedural signals, and/or settings and values (such as resistance, impedance, capacitance, inductance, current/voltage ratings, etc.), that may be applied to influence the operation of a device. Magnetic media, electronic circuits, electrical and optical memory (both volatile and nonvolatile), and firmware are examples of logic. Logic specifically excludes pure signals or software per se (however does not exclude machine memories comprising software and thereby forming configurations of matter).

(146) Within this disclosure, different entities (which may variously be referred to as “units,” “circuits,” other components, etc.) may be described or claimed as “configured” to perform one or more tasks or operations. This formulation—[entity] configured to [perform one or more tasks]—is used herein to refer to structure (i.e., something physical, such as an electronic circuit). More specifically, this formulation is used to indicate that this structure is arranged to perform the one or more tasks during operation. A structure can be said to be “configured to” perform some task even if the structure is not currently being operated. A “credit distribution circuit configured to distribute credits to a plurality of processor cores” is intended to cover, for example, an integrated circuit that has circuitry that performs this function during operation, even if the integrated circuit in question is not currently being used (e.g., a power supply is not connected to it). Thus, an entity described or recited as “configured to” perform some task refers to something physical, such as a device, circuit, memory storing program instructions executable to implement the task, etc. This phrase is not used herein to refer to something intangible.

(147) The term “configured to” is not intended to mean “configurable to.” An unprogrammed FPGA, for example, would not be considered to be “configured to” perform some specific function, although it may be “configurable to” perform that function after programming.

(148) Reciting in the appended claims that a structure is “configured to” perform one or more tasks is expressly intended not to invoke 35 U.S.C. § 112 (f) for that claim element. Accordingly, claims in this application that do not otherwise include the “means for” [performing a function] construct should not be interpreted under 35 U.S.C § 112 (f).

(149) As used herein, the term “based on” is used to describe one or more factors that affect a determination. This term does not foreclose the possibility that additional factors may affect the determination. That is, a determination may be solely based on specified factors or based on the specified factors as well as other, unspecified factors. Consider the phrase “determine A based on B.” This phrase specifies that B is a factor that is used to determine A or that affects the determination of A. This phrase does not foreclose that the determination of A may also be based on some other factor, such as C. This phrase is also intended to cover an embodiment in which A is determined based solely on B. As used herein, the phrase “based on” is synonymous with the phrase “based at least in part on.”

(150) As used herein, the phrase “in response to” describes one or more factors that trigger an effect. This phrase does not foreclose the possibility that additional factors may affect or otherwise trigger the effect. That is, an effect may be solely in response to those factors, or may be in response to the specified factors as well as other, unspecified factors. Consider the phrase “perform A in response to B.” This phrase specifies that B is a factor that triggers the performance of A. This

phrase does not foreclose that performing A may also be in response to some other factor, such as C. This phrase is also intended to cover an embodiment in which A is performed solely in response to B.

(151) As used herein, the terms “first,” “second,” etc. are used as labels for nouns that they precede, and do not imply any type of ordering (e.g., spatial, temporal, logical, etc.), unless stated otherwise. For example, in a register file having eight registers, the terms “first register” and “second register” can be used to refer to any two of the eight registers, and not, for example, just logical registers 0 and 1.

(152) When used in the claims, the term “or” is used as an inclusive or and not as an exclusive or. For example, the phrase “at least one of x, y, or z” means any one of x, y, and z, as well as any combination thereof.

(153) Having thus described illustrative embodiments in detail, it will be apparent that modifications and variations are possible without departing from the scope of the invention as claimed.

Claims

1. A data processor comprising: a plurality of processing elements, each comprising: a weight buffer; an activation buffer; an accumulation memory buffer; a plurality of vector multiply-accumulate units configured to compute in parallel a convolution of weights from the weight buffer and activations from the activation buffer, each of the vector multiply-accumulate units comprising: a first collector disposed between the vector multiply-accumulate unit and the weight buffer; and a second collector disposed between the vector multiply-accumulate unit and the accumulation memory buffer; and configuration logic to configure a depth of the collectors to adjust a level of data-stationary computation of the convolution.
2. The data processor of claim 1, further comprising: the configuration logic to configure a depth of the first collector on one or more of the vector multiply-accumulate units to adjust a level of weight-stationary computation of the convolution.
3. The data processor of claim 1, further comprising: the configuration logic to configure a depth of the second collector on one or more of the vector multiply-accumulate units to adjust a level of output-stationary computation of the convolution.
4. The data processor of claim 1, further comprising: the configuration logic to configure one or both of a depth of the first collector and a depth of the second collector on one or more of the vector multiply-accumulate units to implement one or more of multi-level weight-stationary computation and output-stationary computation of the convolution.
5. The data processor of claim 1, the processing elements further comprising: a third collector disposed between the activation buffer and the vector multiply-accumulate units.
6. The data processor of claim 5, further comprising: the configuration logic to configure a depth of the third collector to adjust a level of input-stationary computation of the convolution.
7. The data processor of claim 5, further comprising: the configuration logic to configure one or more of a depth of the first collector, a depth of the second collector, and a depth of the third collector of one or more of the vector multiply-accumulate units to implement one or more of multi-level weight-stationary, output-stationary, and input stationary computation of the convolution.
8. The data processor of claim 1, further comprising: a global memory; buffer; and the configuration logic to configure the processing elements to utilize the global memory buffer to store the activations and to apply the activations between layers of a neural network.
9. The data processor of claim 1, further comprising: the configuration logic to configure one or more of the vector multiply-accumulate units to compute a portion of the convolution as a partial result and to forward the partial result from the accumulation memory buffer to neighboring

processing elements.

10. The data processor of claim 1, further comprising: the configuration logic to distribute the weights and the activations among the processing elements spatially by a depth of an input of a neural network, and temporally by a height and a width of the input to the neural network.

11. The data processor of claim 1, further comprising: the configuration logic to distribute the weights and the activations among the processing elements spatially and temporally by configurable combinations of input dimensions of a neural network and dimensions of the weights.

12. A neural network computation method comprising: distributing weight values and activation values for a neural network computation among a plurality of processing elements of the neural network spatially by a depth of an input to the neural network, and temporally by a height and a width of the input to the neural network; configuring a depth of at least one collector of a plurality of vector multiply-accumulate units of the processing elements to implement a stationary data flow by the vector multiply-accumulate units during the neural network computation.

13. The neural network computation method of claim 12, wherein the stationary data flow is a weigh-stationary data flow.

14. The neural network computation method of claim 12, wherein the data flow is an output-stationary data flow.

15. A neural network computation method comprising: distributing weight values and activation values for a neural network computation among a plurality of processing elements of the neural network spatially and temporally by configurable combinations of different dimensions of inputs to the neural network and weights of the neural network; and configuring a depth of at least one collector of a plurality of vector multiply-accumulate units of a plurality of processing elements to implement a stationary data flow by the vector multiply-accumulate units during the neural network computation.
