

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12387411
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Hefny; Tarek et al.

Puppeteering a remote avatar by facial expressions

Abstract

A method includes receiving a first facial framework and a first captured image of a face. The first facial framework corresponds to the face at a first frame and includes a first facial mesh of facial information. The method also includes projecting the first captured image onto the first facial framework and determining a facial texture corresponding to the face based on the projected first captured image. The method also includes receiving a second facial framework at a second frame that includes a second facial mesh of facial information and updating the facial texture based on the received second facial framework. The method also includes displaying the updated facial texture as a three-dimensional avatar. The three-dimensional avatar corresponds to a virtual representation of the face.

Inventors: Hefny; Tarek (Redmond, WA), Reiter; Nicholas (Mountain View, CA), Young; Brandon (Mountain View, CA), Kandoor; Arun (Santa Clara, CA), Cower; Dillon (Mountain View, CA)

Applicant: Google LLC (Mountain View, CA)

Family ID: 68467043

Assignee: Google LLC (Mountain View, CA)

Appl. No.: 18/391767

Filed: December 21, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240127523 A1	Apr. 18, 2024

Related U.S. Application Data

continuation parent-doc US 18058621 20221123 US 11887235 child-doc US 18391767
continuation parent-doc US 17052161 US 11538211 20221227 WO PCT/US2019/030218

Publication Classification

Int. Cl.: **G06T13/40** (20110101); **G06T7/13** (20170101); **G06T7/73** (20170101); **G06T17/20** (20060101); **G06T19/20** (20110101); H04L67/10 (20220101)

U.S. Cl.:

CPC **G06T13/40** (20130101); **G06T7/13** (20170101); **G06T7/73** (20170101); **G06T17/20** (20130101); **G06T19/20** (20130101); G06T2207/10024 (20130101); G06T2207/30201 (20130101); H04L67/10 (20130101)

Field of Classification Search

CPC: G06T (19/20); G06T (17/20); G06T (13/40)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2013/0201187	12/2012	Tong et al.	N/A	N/A
2015/0042743	12/2014	Cullen	N/A	N/A
2015/0084950	12/2014	Li et al.	N/A	N/A
2016/0127641	12/2015	Gove	N/A	N/A
2017/0039750	12/2016	Tong et al.	N/A	N/A
2017/0132828	12/2016	Zelenin	N/A	A63F 13/63
2018/0158246	12/2017	Grau	N/A	G06T 3/18

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2011210118	12/2010	JP	N/A
2013020248	12/2012	WO	N/A

OTHER PUBLICATIONS

International Search Report for the related Application No. PCT/US2019/030218 dated May 1, 2019, 16 pages. cited by applicant

China Office Action and Search Report for the related Application No. 201980029370.9, dated Jun. 28, 2022, 34 pages. cited by applicant

Primary Examiner: Salvucci; Matthew

Attorney, Agent or Firm: Honigman LLP

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS (1) This U.S. patent application is a continuation of, and claims priority under 35 U.S.C. § 120 from, U.S. patent application Ser. No. 18/058,621, filed Nov. 23, 2022, which is a continuation of U.S. patent application Ser. No. 17/052,161, filed on Oct. 30, 2020, now U.S. Pat. No. 11,538,211, which is a national stage application of, and claims priority under 35 U.S.C. § 371 from PCT/US2019/030218, filed on May 1, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Application 62/667,767, filed on May 7, 2018. The disclosures of these prior applications are considered part of the disclosure of this application and are hereby incorporated by reference in their entireties.

TECHNICAL FIELD

(1) This disclosure relates to puppeteering a remote avatar by facial expressions.

BACKGROUND

(2) As technology has evolved, people have continued to employ technology as a form of communication. For example, technology allowed communication to expand from a simple physical conversation to a remote real-time conversation. Yet with this expansion, remote forms of communication generally lack some ability to capture expressions and emotions involved in a physical conversation. For example, it often proves difficult to decipher an emotional context from an email or a text conversation. To overcome these deficiencies, methods of communication have sought to provide ways to represent emotion and expression. For example, text applications now include a wide range of emojis and animations to express moods, opinions, or simply offer whimsical entertainment. As people increasingly communicate using real-time audio and video connections, there is an increasing demand for ways to reflect a user's personality and character within these communication channels.

SUMMARY

(3) One aspect of the disclosure provides a method for puppeteering a remote avatar. The method includes receiving, at data processing hardware, a first facial framework and a first captured image of a face of a user with a neutral facial expression. The first facial framework corresponds to the face of the user at a first frame and includes a first facial mesh of facial information. The method also includes projecting, by the data processing hardware, the first captured image of the face onto the first facial framework and determining, by the data processing hardware, a facial texture corresponding to the face of the user based on the projected captured image. The method also includes receiving, at the data processing hardware, a second facial framework that corresponds to the face of the user at a second frame. The second facial framework includes a second facial mesh of facial information. The method also includes updating, by the data processing hardware, the facial texture based on the received second facial framework and displaying, by the data processing hardware, the updated facial texture as a three-dimensional avatar. The three-dimensional avatar corresponds to a virtual representation of the face of the user.

(4) Implementations of the disclosure may include one or more of the following optional features. In some implementations, the method also includes: receiving, at the data processing hardware, a second captured image of the face of the user, the second captured image capturing a smile as a facial expression of the user; receiving, at the data processing hardware, a third captured image of the face of the user, the third captured image capturing, as the facial expression of the user, both eyebrows raised; receiving, at the data processing hardware, a fourth captured image of the face of the user, the fourth captured image capturing, as the facial expression of the user, a smile and both eyebrows raised; for each captured image, determining, by the data processing hardware, a facial expression texture corresponding to the face of the user; blending, by the data processing hardware, the facial expression textures of each captured image and the updated facial texture based on the received second facial framework to generate a blended facial texture; and rendering, by the data processing hardware, the three-dimensional avatar with the blended facial texture. In these implementations, blending further includes: determining a texture vector for each captured image,

the texture vector corresponding to a vector representation of a difference from the first captured image with the neutral facial expression; determining a current texture vector based on the received second facial framework; assigning rendering weights based on a difference between the current texture vector and the texture vector of each captured image; and rendering the three-dimensional avatar with the blended facial texture based on the rendering weights. The rendering weights may have a sum equal to one. In some examples, each of the current texture vector and the texture vector of each captured image may correspond to a fifty-two variable float vector. In these examples, the rendering weights descend in magnitude as the difference between the current texture vector and the texture vector of each captured image increases.

(5) In some examples, the method also includes receiving, at the data processing hardware, a captured current image of the face of the user with a current facial expression mesh of facial information at the second frame, and updating, by the data processing hardware, the facial texture based on the received facial framework and the captured current image. In some implementations, the received captured current image corresponds to a reduced amount of facial texture. In these implementations, the method may also include: determining, by the data processing hardware, an obstructed portion of the face of the user based on the received captured current image; and blending, by the data processing hardware, the obstructed portion of the face of the user with facial texture generated from an unobstructed captured image from a prior frame.

(6) In some implementations, the method also includes generating, by the data processing hardware, a rendition of an eye or a mouth of the user by: detecting, by the data processing hardware, edges of the eye or the mouth; determining, by the data processing hardware, that a sum of angles associated with the edges of the eye or the mouth correspond to two pi (three hundred-sixty degrees); approximating, by the data processing hardware, a position of the eye or the mouth based on the detected edges that correspond to two pi; extracting, by the data processing hardware, the mouth or the eye at the approximated position from the captured image of the face; and rendering, by the data processing hardware, the extracted mouth or the extracted eye at the approximated position with a fill. The captured image may include a red-green-and blue (RGB) image from a mobile phone. The three-dimensional avatar may be displayed on an augmented reality (AR) device.

(7) Another aspect of the disclosure provides a system for puppeteering a remote avatar. The system includes data processing hardware and memory hardware in communication with the data processing hardware, the memory hardware storing instructions that when executed on the data processing hardware cause the data processing hardware to perform operations that include receiving a first facial framework and a first captured image of a face of a user with a neutral facial expression. The first facial framework corresponds to the face of the user at a first frame and includes a first facial mesh of facial information. The operations also include projecting the first captured image of the face onto the first facial framework and determining a facial texture corresponding to the face of the user based on the projected captured image. The operations also include receiving a second facial framework that corresponds to the face of the user at a second frame. The second facial framework includes a second facial mesh of facial information. The operations also include updating the facial texture based on the received second facial framework and displaying the updated facial texture as a three-dimensional avatar. The three-dimensional avatar corresponds to a virtual representation of the face of the user.

(8) This aspect may include one or more of the following optional features. In some implementations, the operations also include: receiving a second captured image of the face of the user, the second captured image capturing a smile as a facial expression of the user; receiving a third captured image of the face of the user, the third captured image capturing, as the facial expression of the user, both eyebrows raised; receiving a fourth captured image of the face of the user, the fourth captured image capturing, as the facial expression of the user, a smile and both eyebrows raised; for each captured image, determining a facial expression texture corresponding to

the face of the user; blending the facial expression textures of each captured image and the updated facial texture based on the received second facial framework to generate a blended facial texture; and rendering the three-dimensional avatar with the blended facial texture. In these implementations, blending further includes: determining a texture vector for each captured image, the texture vector corresponding to a vector representation of a difference from the first captured image with the neutral facial expression; determining a current texture vector based on the received second facial framework; assigning rendering weights based on a difference between the current texture vector and the texture vector of each captured image; and rendering the three-dimensional avatar with the blended facial texture based on the rendering weights. The rendering weights may have a sum equal to one. In some examples, each of the current texture vector and the texture vector of each captured image may correspond to a fifty-two variable float vector. In these examples, the rendering weights descend in magnitude as the difference between the current texture vector and the texture vector of each captured image increases.

(9) In some examples, the operations also include receiving a captured current image of the face of the user with a current facial expression mesh of facial information at the second frame, and updating the facial texture based on the received facial framework and the captured current image. In some implementations, the received captured current image corresponds to a reduced amount of facial texture. In these implementations, the operations may also include: determining an obstructed portion of the face of the user based on the received captured current image; and blending the obstructed portion of the face of the user with facial texture generated from an unobstructed captured image from a prior frame.

(10) In some implementations, the operations also include generating a rendition of an eye or a mouth of the user by: detecting edges of the eye or the mouth; determining that a sum of angles associated with the edges of the eye or the mouth correspond to two pi (three hundred-sixty degrees); approximating a position of the eye or the mouth based on the detected edges that correspond to two pi; extracting the mouth or the eye at the approximated position from the captured image of the face; and rendering the extracted mouth or the extracted eye at the approximated position with a fill. The captured image may include a red-green-and blue (RGB) image from a mobile phone. The three-dimensional avatar may be displayed on an augmented reality (AR) device.

(11) Another aspect of the disclosure provides a method for puppeteering a remote avatar that includes receiving, at data processing hardware, a first facial framework and a first captured image of a face of a user with a neutral facial expression. The first facial framework corresponds to the face of the user at a first frame and includes a first facial mesh of facial information. The method also includes projecting, by the data processing hardware, the first captured image of the face onto the first facial framework and determining, by the data processing hardware, a facial texture corresponding to the face of the user based on the projected first captured image. The method also includes displaying, by the data processing hardware, the determined facial texture as a three-dimensional avatar, the three-dimensional avatar corresponding to a virtual representation of the face of the user.

(12) The details of one or more implementations of the disclosure are set forth in the accompanying drawings and the description below. Other aspects, features, and advantages will be apparent from the description and drawings, and from the claims.

Description

DESCRIPTION OF DRAWINGS

(1) FIG. 1 is a schematic view of an example avatar puppeteering environment.

(2) FIGS. 2A-2F are schematic views of example puppeteers for puppeteering an avatar in the

avatar puppeteering environment of FIG. 1.

(3) FIG. 3 is a flowchart of an example arrangement of operations for a method of puppeteering a remote avatar by facial expressions.

(4) FIG. 4 is a schematic view of an example computing device that may be used to implement the systems and methods described herein.

(5) Like reference symbols in the various drawings indicate like elements.

DETAILED DESCRIPTION

(6) FIG. 1 is an example avatar puppeteering environment **100**. The avatar puppeteering environment **100** is an environment where users **10**, **10a-b** have a conversation **12** via user devices **110**, **110a-b** across a network **120**. The network **120** includes any type of communication network (e.g., a packet switched network) configured to route data between addresses associated with the user devices **110**.

(7) A conversation **12** generally refers to an audible sequence of speech between at least two users **10a**, **10b**. The user device **110** associated with each user **10** is configured to capture and communicate the conversation **12** over the network **120**. The user devices **110** capture not only audio of the speech of the conversation **12**, but also capture images **130** and facial information **140** of faces **20** of the users **10** as the users **10** speak during the conversation. Based on the captured images **130** and the facial information **140** of the faces **20**, each user device **110** are further configured to generate facial expressions **22** for the associated user **10**. Accordingly, the user devices **110** enable remote users **10** to be connected and engaged in a real-time conversation **12**.

(8) The user device **110** can be any computing devices or data processing hardware capable of: (1) communicating facial images **130** and facial information **140** to a network **120** and/or remote system **150**; and (2) displaying a three-dimensional (3D) avatar **160** (e.g., with augmented reality (AR) capabilities). In some examples, a first user device **110a** associated with a first user **10a** is configured to communicate the facial image(s) **130** and facial information **140** associated with the first user **10a**; while a second user device **110b** associated with a second user **10b** is configured to display the 3D avatar **160** associated with the first user **10a**. In the example shown, each user device **110** includes data processing hardware **112**, memory hardware **114**, and one or more image capturing devices **116**. Some examples of image capturing devices **116** are cameras (e.g., depth cameras or RGB cameras) or image sensors (e.g., laser imaging sensors). The user device **110**, includes, but is not limited to, augmented reality (AR) devices, desktop computing devices, and mobile computing devices, such as laptops, tablets, smart phones, and wearable computing devices (e.g., headsets and/or watches). The user devices **110** are configured to utilize their image capturing devices **116** to allow the remote users **10** to engage in conversations **12** across the network **120**.

(9) With continued reference to FIG. 1, each user device **110** executes (i.e., via the data processing hardware **112**) a real-time communication (RTC) application **30** to enable the first and second users **10a**, **10b** to have a conversation **12** with one another. As the first user **10a** speaks to the second user **10b** during the conversation **12**, the first user device **110a** captures audible speech (i.e., audio) **25**, one or more facial images **130** of the face **20** of the first user **10a**, and/or facial information **140** corresponding to the face **20** of the first user **10a**. Thereafter, in some examples, the first user device **110a** transmits an output **118** to the second user device **110b** that includes the captured audible speech **25**, the one or more facial images **130**, and/or the facial information **140** via corresponding audio A and data D channels Ch, but not a video channel Ch, V. Here, the data channel Ch, D includes a lossy data channel configured to transmit the facial images **130** and/or the facial information **140**, while the audio channel Ch, A is configured to communicate the audio **25**. The audible speech **25** transmitted over the audio channel Ch, A includes a digital representation of the speech spoken by the first user **10a**. In other examples, the first user device **110a** transmits the output **118** to the second user device **110b** that includes the audio **25**, the one or more facial images **130**, and/or the facial information **140** via a corresponding video channel Ch, V to ensure synchronization with the related audio **25** from the conversation **12**. For example, synchronizing

the facial images **130** and/or the facial information **140** with the audio **25** via the video channel **Ch**, **V** may be desirable for large data sets (e.g., from the facial images **130** and the facial information **140**) during real-time conversations to avoid latency issues. Optionally, a configuration of the RTC application **30** dictates communication channels **CH** used by the user devices.

(10) Based on the output **118** transmitted from the first user device **110a**, the second user device **110b** is configured to display the 3D avatar **160** corresponding to the face **20** and facial expressions of the first user **10a** on a display **119** of the second user device **110b**. In the example shown, the RTC application **30** executing on the second user device **110b** facilitates communication with a puppeteer **200** that is configured to generate the avatar **160** of the first user **10a** based on the output **118** and provide the generated avatar **160** to the second user device **110b** for display on the display **119**. The 3D avatar **160** generated by the puppeteer **200** corresponds to a virtual representation of the face **20** of the first user **20a**. The puppeteer **200** generates the 3D avatar **160** as a real-time 3D avatar **160** based on the output **160** from the first user device **110a**. In some implementations, the second user device **110b** receives the output **118** including the captured audible speech **25**, the one or more facial images **130**, and/or the facial information **140** from the first user device **110a** via the network **120** and provides the output **118** to the puppeteer **200**. In other implementations, the first user device **110a** transmits the output **118** directly to the puppeteer **200**. In these implementations, the RTC application **30** executing on the first user device **110a** may activate a corresponding 3D avatar feature to allow the first user device **110a** to provide the output **118** directly to the puppeteer **200** for generating the 3D avatar **160** corresponding to the face **20** and facial expressions of the first user **10a**.

(11) In some implementations, the puppeteer **200** includes an application hosted by a remote system **150**, such as a distributed system of a cloud environment, accessed via a user device **110**. In other implementations, the puppeteer **200** includes an application downloaded to memory hardware **114** of the user device **110**. The puppeteer **200** may be configured to communicate with the remote system **150** to access resources **152** (e.g., data processing hardware **154** or memory hardware **156**) for generating the 3D avatar **160** from the facial images **130** and/or the facial information **140**. Additionally or alternatively, the puppeteer **200** may store generated 3D avatars **300** locally on the memory hardware **114** of the user device **110** and/or on the memory hardware **156** of the remote system **150**. For example, the puppeteer **200** and/or the user device **110** may later augment or further render a stored 3D avatar **160** based on later received facial images **130** and/or facial information **140**. Optionally, the RTC application **30** executing on the user device **110** may execute the puppeteer **200** locally without requiring access to the resources **152** of the remote system **150**.

(12) Each facial image **130** refers to an image of the face **20** of a user **10** captured by the image capturing device(s) **116**. The captured facial image **130** may vary in both resolution and embedded data depending on a type of image capturing device **116** that captures the facial image **130**. For example, when a camera or a sensor with depth capability captures the facial image **130** of the user **10**, the captured image **130** includes depth data identifying relationships between facial features and/or facial textures (e.g., shadows, lighting, skin texture, etc.). With depth data, a captured image **130** may inherently include facial information **140** to form a facial mesh **142**. For example, some depth cameras or sensors are configured to generate a mesh from a captured image **130** using surface reconstruction algorithms. In other examples, the captured image **130** generated by cameras or sensors without depth capabilities (e.g., RGB cameras) requires further analysis with techniques such as facial landmark detection and/or facial feature detection to generate facial information **140**.

(13) Facial information **140** generally refers to a point cloud of data related to a face **20**. With facial information **140**, a surface reconstruction algorithm may generate a facial mesh **142** corresponding to the facial information **140**. In some examples, a combination of the facial information **140** and the facial mesh **142** is referred to as a facial framework **144** as this combination corresponds to a facial structure with boundaries associated with the facial information **140**. Although a facial framework **144** bears a resemblance to a user **10**, a facial mesh **142** is generally a smooth rendering

of the facial information **140**. In other words, some unique characteristics of a face **20** of a user **10**, such as wrinkles, dimples, smooth skin, dry skin, oily skin, porosity, etc., are lost with translation of the user **10** to a facial framework **144**. To account for these missing aspects, the puppeteer **200** is configured to generate a facial texture **212** corresponding to these unique characteristics based on the facial framework **144** and at least one captured image **130**.

(14) FIGS. 2A-2F are examples of the puppeteer **200** generating the 3D avatar **160** based on the received output **118** including the captured image(s) **130** and the facial information **140**. The puppeteer **200** includes a texturer **210** and an updater **220**. The texturer **210** is configured to determine a facial texture **212**, while the updater **220** is configured to update the facial texture **212** based on subsequently received facial framework(s) **144** and/or captured image(s) **130**. Referring to FIG. 2A, the puppeteer **200** receives the output **118** corresponding to a first captured image **130** of the face **20** of the user **10** and a first facial framework **144a** that includes a first facial mesh **142** of facial information **140** of the user **10**. The first facial framework **144a** corresponds to the face **20** of the user **10** at a first frame F1. The captured image **130** may capture a facial expression **22** of the user **10**. For instance, the captured image **130** may include a neutral facial expression **22a** of the user **10**. In the example shown, the texturer **210** projects the first captured image **130** of the face **20** onto the first facial framework **144a** to determine a facial texture **212**, **212a** corresponding to the neutral facial expression **22**, **22a** of the face **20**. After the texturer **210** determines the facial texture **212** (e.g., the first facial texture **212a**), the updater **220** may then update the facial texture **212** based on a subsequent frame F2 from the conversation **12** that occurs subsequent in time to the first frame F1 to form an updated facial texture **212**, **212U**. Based on this updated facial texture **212U**, the puppeteer **200** updates the displayed 3D avatar **160** as a user **10** proceeds to talk and change facial expressions **22** in real-time.

(15) Implementations include the puppeteer **200** operating with minimal bandwidth requirements. Being conscientious of bandwidth, the facial texture **212** determined by the texturer **210** includes a static texture that updates based solely on facial information **140**, such as facial framework(s) **144**. In other words, rather relying on large captured image files at the puppeteer **200**, the puppeteer **200** generates the 3D avatar **160** by updating the facial texture **212** based on facial information of later frames F in the conversation **12** (e.g., the second frame F2). This static approach permits updates to the facial mesh **142** and the facial structure **144** in real-time without incurring increased bandwidth requirements of the avatar puppeteering environment **100**. For example, FIG. 2A shows the updater **220** of the puppeteer **200** receiving the facial texture **212a** in combination with a second facial framework **144b** corresponding to the face **20** of the user **10** at a second frame F2. Here, much like the first frame F1, the second facial framework **144b** includes a second facial mesh **142b** of the facial information **140** at the second frame F2. In this configuration, the updater **220** updates the facial texture **212**, **212a** based on the received second facial framework **144b** to form the updated facial texture **212U**. Thereafter, the puppeteer **200** uses the updated facial texture **212U** to generate the 3D avatar **160** (or update an existing 3D avatar **160**) and provides the generated 3D avatar **160** to the user device **110** for display on the display **119**. In some examples, the puppeteer **200** provides the updated facial texture **212U** to the user device **110** and the user device **110** (e.g., via the RTC application **30**) generates the 3D avatar **160** or updates an existing 3D avatar **160**.

(16) Referring to FIG. 2B, in some implementations, the puppeteer **200** receives a plurality of captured images **130**, **130a-d** of the face **20** of the user **20** and determines, for each captured image **130**, a corresponding facial texture **212**, **212a-d** by projecting the captured image **130** of the face **20** onto the first facial framework **140a**. Thereafter, the puppeteer **200** updates and blends each facial texture **212a-d** based on the second facial framework **140b** to generate a blended facial texture **212.sub.blend**. While FIG. 2B depicts using four captured images **130a-d** to generate the blended facial texture **212.sub.blend**, any number of captured images **130** may be used without departing from the scope of the present disclosure. As such, by incorporating more than one captured image **130** into the facial texture generation, the puppeteer **200** may account for other baseline facial

expressions **22**, **22a-d** in addition to, or in lieu of, the single neutral facial expression **22a** of FIG. 2A.

(17) In the example shown, the puppeteer **200** receives the output **118** corresponding to four captured images **130**, **130a-d** of the face **20** of the user **10** and the first facial framework **144a** that includes the first facial mesh **142** of the facial information **140** of the user **10** at the first frame **F1**. The first facial framework **144a** corresponds to the face **20** of the user **10** at a first frame **F1**. Here, each captured image **130a-d** corresponds to a different facial expression **22**, **22a-d** of the user facial expression **22** of the user **10**. For instance, the first captured image **130a** corresponds to the neutral facial expression **22a**, the second captured image **130b** corresponds to a smiling facial expression **22b**, the third captured image **130c** corresponds to a both eyebrows raised facial expression **22c**, and the fourth captured image **130d** corresponds to a smiling with both eyebrows raised facial expression **22d**. Accordingly, the texturer **210** is configured to determine a corresponding facial texture **212**, **212a-d** for each captured image **130** by projecting the captured image **130** onto the first facial framework **144a**.

(18) With continued reference to FIG. 2B, the updater **220** receives the facial textures **212a-d** from the texturer **210**. In some examples, the updater **220** updates each facial texture **212a-d** based on the received second facial framework **144b** and blends the corresponding updated facial textures **212U** together to generate the blended facial texture **212.sub.blend** at the second frame **F2**.

Thereafter, the puppeteer **200** uses the blended facial texture **212.sub.blend** to generate the 3D avatar **160** (or update an existing 3D avatar **160**) and provides the generated 3D avatar **160** to the user device **110** for display on the display **119**. In some examples, the puppeteer **200** provides the blended facial texture **212.sub.blend** to the user device **110** and the user device **110** (e.g., via the RTC application **30**) generates the 3D avatar **160** or updates an existing 3D avatar **160**.

(19) Referring to FIG. 2C, in some examples, the puppeteer **200** further includes vector and weight generators **230**, **240** that cooperate to provide the updater **220** with rendering weights **244** for updating and blending the four facial textures **212a-d** output from the texturer **210** to generate the blended facial texture **212.sub.blend** at a current frame **F** (e.g., the second frame **F2** in FIG. 2C). In the example shown, the vector generator **230** receives each facial texture **212a-d** output from the texturer **210** and generates corresponding texture vectors **232** relative to a baseline facial texture **212**. For instance, the baseline facial texture **212** may correspond to the first facial texture **212a** associated with the first captured image **130a** corresponding to the natural facial expression **22a**. As such, the vector generator **230** may generate a first texture vector **232ba** based on the second facial texture **212b** relative to the first facial texture **212a**, a second texture vector **232ca** based on the third facial texture **212c** relative to the first facial texture **212a**, and a third texture vector **232da** based on the fourth facial texture **212d** relative to the first facial texture **212a**. Further, the vector generator **230** generates a current texture vector **232**, **232current** corresponding to the facial information **140** at a recent frame **F** (e.g., the second frame **F2**). For instance, the vector generator **230** generates the current texture vector **232current** between the first facial framework **144a** at the first frame **F1** and the second facial framework **144b** at the second frame **F2**.

(20) The weight generator **240** receives the current texture vector **232.sub.current** and each of the texture vectors **232ba-da** from the vector generator **230** and generates rendering weights **244** based on a respective difference **242** between the current texture vector **232.sub.current** and each texture vector **232ba-da**. In other words, the rendering weights **244** account for deviations at a current frame relative to the facial textures **212**, **212a-d**. Rendering weights **244** may be configured to correspond to known detected facial expressions. For example, the rendering weights **244** may include vectors associated with locations of facial landmarks such that each vector represents a magnitude and a direction from a baseline position of a facial landmark (e.g., from the first facial framework **144a** of a neutral facial expression **22**, **22a** to a second facial framework **144b** of a facial expression **22** at the second frame **F2**). In one example, the rendering weights **244** form a fifty-two variable float vector. In some examples, the rendering weights **244** correspond to blending

percentages such that values of the rendering weights **244** include respective ratios with a sum equal to one.

(21) In some implementations, the weight generator **240** assigns a highest value to a value within the rendering weights **244** when a texture vector **232** is closest to the current texture vector **232.sub.current**. For example, when the second facial framework **144b** indicates that a current facial expression **22** of the user **10** approaches a smile (e.g., the second facial expression **22b** associated with the second texture vector **232ba**), the respective difference between the current texture vector **232.sub.current** and the second texture vector **232ba** associated with the smiling facial expression **22b** is less than the differences between the current texture vector **232.sub.current** and the other texture vectors **232ca**, **232da**. In this instance, the weight generator **240** assigns values to the rendering weights **244** bias toward the smiling facial expression **22b** (e.g., a higher rendering weight value). Accordingly, updater **220** uses these rendering weights **244** assigned by the weight generator **240** generate the blended facial texture **222.sub.blend** more towards the second facial texture **212b** associated with the smiling facial expression **22b**.

(22) Unlike the puppeteer **200** of FIG. 2A operating with minimal bandwidth requirements, the puppeteers **200** of FIGS. 2B and 2C require greater bandwidth by accounting for more captured images **130a-d** of the face **20** of the user **10** to achieve a more accurate visual representation of the face **20** of the user **10**. Here, a puppeteer **200** with a finite number of captured images **130** (e.g., four captured images **130a-d**) may increase accuracy while still minimizing bandwidth by updating the 3D avatar **160** based on facial information **140** (e.g., a second facial framework **144b**) at a current frame (e.g., the second frame F2) rather than updating the facial texture **212** from a current captured image **130**, **130.sub.current** (as shown in FIG. 2D).

(23) FIG. 2D is an examples of a puppeteer **200** that receives the current captured image **130**, **130.sub.current** at the second frame F2. In this configuration, the puppeteer **200** operates similar to the puppeteer **200** of FIG. 2A except that the updater **220** updates the first facial texture **212a** based on both the second facial framework **144b** and the current captured image **130.sub.current**. In some implementations, when utilizing the current captured image **130.sub.current** of the user **10**, the puppeteer **200** receives and/or reduces an amount of facial texture **212** associated with the current captured image **130.sub.current**. For example, the updater **220** generates the updated facial texture **212U** based on the current captured image **130.sub.current** having one third of the facial texture **212** (e.g., when compared to the first facial texture **212a**). By reducing an amount of facial texture **212** within the current captured image **130.sub.current**, the puppeteer **200** may reduce its operating bandwidth requirements.

(24) Referring to FIG. 2E, in some examples, facial information **140** and/or facial framework(s) **144** correspond to a partial capture (e.g., an obstructed image **214**) of the face **20** of the user **10**. For example, the user **10** moves within a field of view or moves the image capturing device **116**. In these examples, the puppeteer **200** may be additionally configured to account for these issues. In some configurations, the texturer **210** identifies whether the current capture image **130.sub.current** and or second facial framework **144b** corresponds to an obstructed image. For example, the texturer **210** tracks and analyzes how much facial information **140** is received on average and compares this data to the current capture image **130.sub.current** and/or the second facial framework **144b**. When the texturer **210** identifies an obstructed image and/or obstructed facial information, the texturer **210** identifies a preceding frame F.sub.n-1 that is not obstructed to generate the facial texture **212** for the obstructed portion of the obstructed capture **214**. For example, when the texturer **210** determines the second frame F.sub.2 includes an obstructed image **214** and the first frame F1 includes an unobstructed image (e.g., the first captured image **130a**), the texturer **210** may render the obstructed capture **214** (e.g., the received current captured image **130.sub.current**) with the facial information **140** associated with the first frame F.sub.1.

(25) Referring to FIG. 2F, in some implementations, the puppeteer **200** includes a feature filler **250**. The feature filler **250** identifies often troublesome features like eyes or mouths and fills in (i.e.

visually represents) cavities associated with these features. FIG. 2F shows a simplified puppeteer **200** to focus on the feature filler **250**. In some examples, the feature filler **250** detects edges of features. For example, the feature filler **250** sums all angles that center around a vertex. When the sum equals two pi or three hundred and sixty degrees, the feature filler **250** determines that the feature is a cavity, such as an eye or mouth. When the sum does not equal two pi, the feature filler **250** identifies the feature as an edge vertex. Once the feature is identified as a cavity, the feature filler **250** approximates a position of the cavity based on facial proportions and/or locations of the detected edges. Here, at the approximated position, the feature filler **250** extracts the feature and renders the extracted feature with a fill. In some examples, a two-ear approach is used to fill the feature while the facial texture **212** maps vertices used during edge detection for the feature filler **250**.

(26) FIG. 3 is a flowchart for an example arrangement of operations for a method **300** of puppeteering a remote avatar **160**. At operation **302**, the method **300** receives a first facial framework **144**, **144a** and a first captured image **130** of a face **20** of the user **10** with a neutral facial expression **22**, **22a**. The first facial framework **144a** corresponds to the face **20** of the user **10** at a first frame F1 and includes a first facial mesh **142**, **142a** of facial information **140**. At operation **304**, the method **300** projects the first captured image **130** of the face **20** onto the first facial framework **144a**. At operation **306**, the method **300** determines a facial texture **212** corresponding to the face **20** of the user **10** based on the projected captured image **130**. At operation **308**, the method **300** receives a second facial framework **144b** corresponding to the face **20** of the user **10** at a second frame F2. The second facial framework **144b** includes a second facial mesh **142**, **142b** of the facial information **140**. At operation **310**, the method **300** updates the facial texture **212** based on the received second facial framework **144b**. At operation **312**, the method **300** displays the updated facial texture **212** as a 3D avatar **160**. The 3D avatar **160** corresponds to a virtual representation of the face **20** of the user **10**.

(27) A software application (i.e., a software resource) may refer to computer software that causes a computing device to perform a task. In some examples, a software application may be referred to as an “application,” an “app,” or a “program.” Example applications include, but are not limited to, system diagnostic applications, system management applications, system maintenance applications, word processing applications, spreadsheet applications, messaging applications, media streaming applications, social networking applications, and gaming applications.

(28) FIG. 4 is schematic view of an example computing device **400** that may be used to implement the systems and methods of, for example, the user device **110**, the remote system **150**, and the puppeteer **200**, described in this document. The computing device **400** is intended to represent various forms of digital computers, such as laptops, desktops, workstations, personal digital assistants, servers, blade servers, mainframes, and other appropriate computers. The components shown here, their connections and relationships, and their functions, are meant to be exemplary only, and are not meant to limit implementations of the inventions described and/or claimed in this document.

(29) The computing device **400** includes a processor **410**, memory **420**, a storage device **430**, a high-speed interface/controller **440** connecting to the memory **420** and high-speed expansion ports **450**, and a low speed interface/controller **460** connecting to a low speed bus **470** and a storage device **430**. Each of the components **410**, **420**, **430**, **440**, **450**, and **460**, are interconnected using various busses, and may be mounted on a common motherboard or in other manners as appropriate. The processor **410** can process instructions for execution within the computing device **400**, including instructions stored in the memory **420** or on the storage device **430** to display graphical information for a graphical user interface (GUI) on an external input/output device, such as display **480** coupled to high speed interface **440**. In other implementations, multiple processors and/or multiple buses may be used, as appropriate, along with multiple memories and types of memory. Also, multiple computing devices **400** may be connected, with each device providing portions of

the necessary operations (e.g., as a server bank, a group of blade servers, or a multi-processor system).

(30) The memory **420** stores information non-transitorily within the computing device **400**. The memory **420** may be a computer-readable medium, a volatile memory unit(s), or non-volatile memory unit(s). The non-transitory memory **420** may be physical devices used to store programs (e.g., sequences of instructions) or data (e.g., program state information) on a temporary or permanent basis for use by the computing device **400**. Examples of non-volatile memory include, but are not limited to, flash memory and read-only memory (ROM)/programmable read-only memory (PROM)/erasable programmable read-only memory (EPROM)/electronically erasable programmable read-only memory (EEPROM) (e.g., typically used for firmware, such as boot programs). Examples of volatile memory include, but are not limited to, random access memory (RAM), dynamic random access memory (DRAM), static random access memory (SRAM), phase change memory (PCM) as well as disks or tapes.

(31) The storage device **430** is capable of providing mass storage for the computing device **400**. In some implementations, the storage device **430** is a computer-readable medium. In various different implementations, the storage device **430** may be a floppy disk device, a hard disk device, an optical disk device, or a tape device, a flash memory or other similar solid state memory device, or an array of devices, including devices in a storage area network or other configurations. In additional implementations, a computer program product is tangibly embodied in an information carrier. The computer program product contains instructions that, when executed, perform one or more methods, such as those described above. The information carrier is a computer- or machine-readable medium, such as the memory **420**, the storage device **430**, or memory on processor **410**.

(32) The high speed controller **440** manages bandwidth-intensive operations for the computing device **400**, while the low speed controller **460** manages lower bandwidth-intensive operations. Such allocation of duties is exemplary only. In some implementations, the high-speed controller **440** is coupled to the memory **420**, the display **480** (e.g., through a graphics processor or accelerator), and to the high-speed expansion ports **450**, which may accept various expansion cards (not shown). In some implementations, the low-speed controller **460** is coupled to the storage device **430** and a low-speed expansion port **490**. The low-speed expansion port **490**, which may include various communication ports (e.g., USB, Bluetooth, Ethernet, wireless Ethernet), may be coupled to one or more input/output devices, such as a keyboard, a pointing device, a scanner, or a networking device such as a switch or router, e.g., through a network adapter.

(33) The computing device **400** may be implemented in a number of different forms, as shown in the figure. For example, it may be implemented as a standard server **400a** or multiple times in a group of such servers **400a**, as a laptop computer **400b**, or as part of a rack server system **400c**.

(34) Various implementations of the systems and techniques described herein can be realized in digital electronic and/or optical circuitry, integrated circuitry, specially designed ASICs (application specific integrated circuits), computer hardware, firmware, software, and/or combinations thereof. These various implementations can include implementation in one or more computer programs that are executable and/or interpretable on a programmable system including at least one programmable processor, which may be special or general purpose, coupled to receive data and instructions from, and to transmit data and instructions to, a storage system, at least one input device, and at least one output device.

(35) These computer programs (also known as programs, software, software applications or code) include machine instructions for a programmable processor, and can be implemented in a high-level procedural and/or object-oriented programming language, and/or in assembly/machine language. As used herein, the terms “machine-readable medium” and “computer-readable medium” refer to any computer program product, non-transitory computer readable medium, apparatus and/or device (e.g., magnetic discs, optical disks, memory, Programmable Logic Devices (PLDs)) used to provide machine instructions and/or data to a programmable processor, including a

machine-readable medium that receives machine instructions as a machine-readable signal. The term “machine-readable signal” refers to any signal used to provide machine instructions and/or data to a programmable processor.

(36) The processes and logic flows described in this specification can be performed by one or more programmable processors executing one or more computer programs to perform functions by operating on input data and generating output. The processes and logic flows can also be performed by special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit). Processors suitable for the execution of a computer program include, by way of example, both general and special purpose microprocessors, and any one or more processors of any kind of digital computer. Generally, a processor will receive instructions and data from a read only memory or a random access memory or both. The essential elements of a computer are a processor for performing instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto optical disks, or optical disks. However, a computer need not have such devices. Computer readable media suitable for storing computer program instructions and data include all forms of non-volatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto optical disks; and CD ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

(37) To provide for interaction with a user, one or more aspects of the disclosure can be implemented on a computer having a display device, e.g., a CRT (cathode ray tube), LCD (liquid crystal display) monitor, or touch screen for displaying information to the user and optionally a keyboard and a pointing device, e.g., a mouse or a trackball, by which the user can provide input to the computer. Other kinds of devices can be used to provide interaction with a user as well; for example, feedback provided to the user can be any form of sensory feedback, e.g., visual feedback, auditory feedback, or tactile feedback; and input from the user can be received in any form, including acoustic, speech, or tactile input. In addition, a computer can interact with a user by sending documents to and receiving documents from a device that is used by the user; for example, by sending web pages to a web browser on a user's client device in response to requests received from the web browser.

(38) Further to the descriptions above, a user may be provided with controls allowing the user to make an election as to both if and when systems, programs or features described herein may enable collection of user information (e.g., information about a user's social network, social actions or activities, profession, a user's preferences, or a user's current location), and if the user is sent content or communications from a server. In addition, certain data may be treated in one or more ways before it is stored or used, so that personally identifiable information is removed. For example, a user's identity may be treated so that no personally identifiable information can be determined for the user, or a user's geographic location may be generalized where location information is obtained (such as to a city, ZIP code, or state level), so that a particular location of a user cannot be determined. Thus, the user may have control over what information is collected about the user, how that information is used, and what information is provided to the user.

(39) A number of implementations have been described. Nevertheless, it will be understood that various modifications may be made without departing from the spirit and scope of the disclosure. Accordingly, other implementations are within the scope of the following claims.

Claims

1. A computer-implemented method when executed by data processing hardware causes the data processing hardware to perform operations comprising: receiving a first facial framework and a first captured image of a face of a user with a neutral facial expression, the first facial framework corresponding to the face of the user at a first frame; determining a first facial texture based on the first facial framework and the first captured image, the first facial texture corresponding to the neutral facial expression; receiving a second facial framework and a second captured image of the face of the user with a non-neutral facial expression, the second facial framework corresponding to the face of the user at a second frame; determining a second facial texture based on the second facial framework, the second facial texture corresponding to the non-neutral facial expression; blending the neutral facial expression corresponding to the first facial texture and the non-neutral facial expression corresponding to the second facial texture together to generate a third facial texture; and displaying the third facial texture as a three-dimensional avatar, the three-dimensional avatar corresponding to a virtual representation of the face of the user.
2. The computer-implemented method of claim 1, wherein the operations further comprise: receiving audio data associated with the first captured image; and generating the audio data synchronously with displaying the third facial texture as the three-dimensional avatar.
3. The computer-implemented method of claim 1, wherein receiving the first captured image comprises receiving the first captured image from a first user device associated with the user.
4. The computer-implemented method of claim 3, wherein displaying the third facial texture as the three-dimensional avatar comprises displaying the third facial texture as the three-dimensional avatar at a second user device associated with another user.
5. The computer-implemented method of claim 1, wherein the non-neutral facial expression comprises a smiling facial expression.
6. The computer-implemented method of claim 1, wherein the non-neutral facial expression comprises both eyebrows raised.
7. The computer-implemented method of claim 1, wherein the first captured image comprises a red-green-and blue image from a first user device.
8. The computer-implemented method of claim 1, wherein the three-dimensional avatar is displayed on an augmented reality device.
9. The computer-implemented method of claim 1, wherein the first facial framework comprises an obstructed view of the face of the user.
10. The computer-implemented method of claim 1, wherein the operations further comprise rendering a facial cavity onto the first facial framework.
11. A system comprising: data processing hardware; and memory hardware in communication with the data processing hardware, the memory hardware storing instructions that when executed on the data processing hardware cause the data processing hardware to perform operations comprising: receiving a first facial framework and a first captured image of a face of a user with a neutral facial expression, the first facial framework corresponding to the face of the user at a first frame; determining a first facial texture based on the first facial framework and the first captured image, the first facial texture corresponding to the neutral facial expression; receiving a second facial framework and a second captured image of the face of the user with a non-neutral facial expression, the second facial framework corresponding to the face of the user at a second frame; determining a second facial texture based on the second facial framework, the second facial texture corresponding to the non-neutral facial expression; blending the neutral facial expression corresponding to the first facial texture and the non-neutral facial expression corresponding to the second facial texture together to generate a third facial texture; and displaying the third facial texture as a three-dimensional avatar, the three-dimensional avatar corresponding to a virtual representation of the face of the user.
12. The system of claim 11, wherein the operations further comprise: receiving audio data

associated with the first captured image; and generating the audio data synchronously with displaying the third facial texture as the three-dimensional avatar.

13. The system of claim 11, wherein receiving the first captured image comprises receiving the first captured image from a first user device associated with the user.

14. The system of claim 13, wherein displaying the third facial texture as the three-dimensional avatar comprises displaying the third facial texture as the three-dimensional avatar at a second user device associated with another user.

15. The system of claim 11, wherein the non-neutral facial expression comprises a smiling facial expression.

16. The system of claim 11, wherein the non-neutral facial expression comprises both eyebrows raised.

17. The system of claim 11, wherein the first captured image comprises a red-green-and blue image from a first user device.

18. The system of claim 11, wherein the three-dimensional avatar is displayed on an augmented reality device.

19. The system of claim 11, wherein the first facial framework comprises an obstructed view of the face of the user.

20. The system of claim 11, wherein the operations further comprise rendering a facial cavity onto the first facial framework.
