

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12387473
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Jason; Adam

Neural network host platform for detecting anomalies in cybersecurity modules

Abstract

Aspects of the disclosure relate to anomaly detection in cybersecurity training modules. A computing platform may receive information defining a training module. The computing platform may capture a plurality of screenshots corresponding to different permutations of the training module. The computing platform may input, into an auto-encoder, the plurality of screenshots corresponding to the different permutations of the training module, wherein inputting the plurality of screenshots corresponding to the different permutations of the training module causes the auto-encoder to output a reconstruction error value. The computing platform may execute an outlier detection algorithm on the reconstruction error value, which may cause the computing platform to identify an outlier permutation of the training module. The computing platform may generate a user interface comprising information identifying the outlier permutation of the training module. The computing platform may send the user interface to at least one user device.

Inventors:	Jason; Adam (Zelienople, PA)
Applicant:	Proofpoint, Inc. (Sunnyvale, CA)
Family ID:	76180938
Assignee:	Proofpoint, Inc. (Sunnyvale, CA)
Appl. No.:	18/374274
Filed:	September 28, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240020960 A1	Jan. 18, 2024

Related U.S. Application Data

Publication Classification

Int. Cl.: **G06V10/778** (20220101); **G06F18/21** (20230101); **G06F18/214** (20230101);
G06F18/2433 (20230101); **G06N3/02** (20060101); **G06V10/94** (20220101); H04L9/40
(20220101)

U.S. Cl.:

CPC **G06V10/7784** (20220101); **G06F18/2148** (20230101); **G06F18/2185** (20230101);
G06F18/2433 (20230101); **G06N3/02** (20130101); **G06V10/95** (20220101); H04L63/14
(20130101)

Field of Classification Search

CPC: G06K (9/6284); G06K (9/6257); G06K (9/6264); G06N (3/02); H04L (63/14)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2005/0108619	12/2004	Theall	715/255	G06F 16/48
2012/0151309	12/2011	Bastide	715/202	G06F 40/186
2013/0083996	12/2012	Prasad et al.	N/A	N/A
2019/0130279	12/2018	Beggel et al.	N/A	N/A
2019/0324890	12/2018	Li et al.	N/A	N/A

OTHER PUBLICATIONS

Sarajanovic-Djukic et al, “Fast Distance-based Anomaly Detection in Images Using an Inception-like Autoencoder”, Mar. 12, 2020, arXiv:2003.08731v1, pp. 1-15 (Year: 2020). cited by examiner
Oct. 27, 2021—(EP) Extended Search Report—App 21176499.8. cited by applicant
Sarajanovic-Djuikic et al., “Fast Distance-based Anomaly Detection in Images Using an Inception-like Autoencoder”, Mar. 12, 2020, arXiv:2003.08731v,, pp. 1-15 (Year: 2020). cited by applicant
Nov. 13, 2024—(EP) Communication under Rule 71(3) EPC—App 21176499.8. cited by applicant

Primary Examiner: Dunphy; David F
Attorney, Agent or Firm: Banner & Witcoff, Ltd.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of and claims priority to co-pending U.S. application Ser. No. 17/740,740, filed May 10, 2022, and entitled, “Neural Network Host Platform for Detecting Anomalies in Cybersecurity Modules,”

which is a continuation of and claims priority to U.S. application Ser. No. 17/038,727, filed Sep. 30, 2020, and entitled, "Neural Network Host Platform for Detecting Anomalies in Cybersecurity Modules," which claims the benefit of and priority to U.S. Provisional Patent Application Ser. No. 63/040,770, filed Jun. 18, 2020, and entitled "Providing Test Automation and Uniformity Analysis," all of which are incorporated by reference herein in their entirety.

TECHNICAL FIELD

(1) Aspects of the disclosure relate to data processing methods, machine learning systems, and communication systems and networks. In particular, one or more aspects of the disclosure relate to identifying anomalies in cybersecurity modules using machine learning.

BACKGROUND

(2) Increasingly, organizations face various cybersecurity threats. Various modules may be developed to provide cybersecurity training to employees of these organizations so as to combat these threats. It remains difficult, however, to automatically achieve uniformity throughout different permutations of these modules. This is particularly true in circumstances where large numbers of permutations exist for each module. Undetected anomalies in these permutations may result in visual defects and/or other deficiencies that may cause poor user experiences. Attempts to integrate such analysis techniques into efficient and effective automated processes present various technical challenges, particularly when trying to balance module uniformity against the optimal consumption of computing resources, such as processing power and network bandwidth.

SUMMARY

(3) Aspects of the disclosure provide technical solutions that overcome one or more of the technical problems described above and/or other technical challenges. For instance, one or more aspects of the disclosure relate to automatically identifying anomalies in cybersecurity training modules using machine learning.

(4) In accordance with one or more embodiments, a computing platform having at least one processor, a communication interface, and memory may receive information defining a training module. The computing platform may capture a plurality of screenshots corresponding to different permutations of the training module. The computing platform may input, into an auto-encoder, the plurality of screenshots corresponding to the different permutations of the training module, which may cause the auto-encoder to output a reconstruction error value. The computing platform may execute, on the reconstruction error value, an outlier detection algorithm, which may cause the computing platform to identify an outlier permutation of the training module. The computing platform may generate a user interface comprising information identifying the outlier permutation of the training module. The computing platform may send the user interface to at least one user device.

(5) In one or more instances, the auto-encoder may be trained on a set of images corresponding to a plurality of training modules including the training module. In one or more instances, the set of images may include images configured for one or more of: a desktop computer, a laptop computer, or a mobile device.

(6) In one or more instances, the user interface may include controls allowing a user of the at least one user device to edit the outlier permutation of the training module. In one or more instances, the different permutations of the training module may correspond to one or more of:

(7) different languages, different browsers, or different resolutions.

(8) In one or more instances, the computing platform may process, prior to inputting the plurality of screenshots into the auto-encoder, the plurality of screenshots, which may include redacting one or more of: text, images, or template components from each of the plurality of screenshots. In one or more instances, the computing platform may use the information defining the training module to render the different permutations of the training module.

(9) In one or more instances, sending the user interface to the at least one user device may cause the

at least one user device to display the user interface. In one or more instances, the user interface may include one or more of: an alpha-numeric identifier for the outlier permutation of the training module, a link to access the outlier permutation of the training module, a summary of corrections to be made to the outlier permutation of the training module, an individual assigned to perform the corrections, or a status of the corrections.

(10) In one or more instances, the reconstruction error value may indicate a degree to which the outlier permutation of the training module differs from an anticipated permutation of the training module. In one or more instances, the computing platform may identify the outlier permutation of the training module by: 1) identifying, using the outlier detection algorithm and the reconstruction error value, that the degree to which the outlier permutation of the training module differs from the anticipated permutation of the training module exceeds a predetermined anomaly identification threshold; and 2) based on the identification that the degree to which the outlier permutation of the training module differs from the anticipated permutation of the training module exceeds the predetermined anomaly identification threshold, identifying the outlier permutation of the training module.

(11) In one or more instances, after sending the user interface to the at least one user device, the computing platform may receive user input indicating whether or not the outlier permutation of the training module was correctly identified as an outlier. Based on the user input indicating whether or not the outlier permutation of the training module was correctly identified as an outlier, the computing platform may dynamically tune the auto-encoder.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present disclosure is illustrated by way of example and not limited in the accompanying figures in which like reference numerals indicate similar elements and in which:

(2) FIG. 1 depicts an illustrative operating environment for applying machine learning in anomaly detection in accordance with one or more example embodiments;

(3) FIGS. 2A-2E depict an illustrative event sequence for applying machine learning in anomaly detection in accordance with one or more example embodiments;

(4) FIGS. 3 and 4 depict illustrative user interfaces for applying machine learning in anomaly detection in accordance with one or more example embodiments; and

(5) FIG. 5 depicts an illustrative method for applying machine learning in anomaly detection in accordance with one or more example embodiments.

DETAILED DESCRIPTION

(6) In the following description of various illustrative embodiments, reference is made to the accompanying drawings, which form a part hereof, and in which is shown, by way of illustration, various embodiments in which aspects of the disclosure may be practiced. It is to be understood that other embodiments may be utilized, and structural and functional modifications may be made, without departing from the scope of the present disclosure. Various connections between elements are discussed in the following description. It is noted that these connections are general and, unless specified otherwise, may be direct or indirect, wired or wireless, and that the specification is not intended to be limiting in this respect.

(7) As a brief introduction to the concepts described further below, one or more aspects of the disclosure relate to systems and methods for automatically analyzing cybersecurity training modules to ensure uniformity. For example, many permutations of various training modules may be supported by an enterprise organization so as to configure each training module for various different languages, screen resolutions, browsers, and/or other parameters. Given all of these permutations, it may be difficult to ensure that all permutations of the modules are visually correct

(e.g., that each permutation involves displaying the corresponding cybersecurity training module as it is intended to be displayed).

(8) Accordingly, described herein is a machine learning tool that learns what the layout of each cybersecurity module should look like, analyzes all corresponding permutations of each cybersecurity module, and flags identified anomalies for review and repair. In doing so, one or more of the systems and methods described herein solve technical problems corresponding to module scalability, overcome challenges (such as maintaining module uniformity) corresponding to visual testing, provide an improvement over alternative approaches to such uniformity analysis such as pixel to pixel comparison, and/or address additional technical deficiencies related to maintaining module uniformity.

(9) Furthermore, one or more of the systems and methods described herein address deficiencies related to manual review of training modules. For example, it might not be feasible for a human to manually perform a regression test to all permutations of a set of training modules. For instance, there may be fifty different modules each produced for a plurality of locales, breakpoints, browsers, or other specified characteristics, and each module may contain tens or hundreds of pages. If such anomalies are not detected, the modules may suffer from visual defects that may result in a poor user experience (e.g., containing unreadable training information or other visual defects). In some instances, due to the volume of modules and variations thereof, locating these visual defects may be a time consuming and/or error prone process. Accordingly, by applying the techniques described herein, the process of anomaly detection may be streamlined for efficiency and increased accuracy, and a fast, low cost method for visual testing of modules may be provided.

(10) FIG. 1 depicts an illustrative operating environment for applying machine learning in anomaly detection in accordance with one or more example embodiments. Referring to FIG. 1, computing environment **100** may include various computer systems, computing devices, networks, and/or other operating infrastructure. For example, computing environment **100** may include a neural network host platform **110**, cybersecurity module host system **120**, enterprise user device **130**, administrator computing device **140**, and a network **190**.

(11) Network **190** may include one or more wired networks and/or one or more wireless networks that interconnect neural network host platform **110**, cybersecurity module host system **120**, enterprise user device **130**, administrator computing device **140**, and/or other computer systems and/or devices. In addition, each of neural network host platform **110**, cybersecurity module host system **120**, enterprise user device **130**, administrator computing device **140** may be special purpose computing devices configured to perform specific functions, as illustrated in greater detail below, and may include specific computing components such as processors, memories, communication interfaces, and/or the like.

(12) Neural network host platform **110** may include one or more processor(s) **111**, one or more memory(s) **112**, and one or more communication interface(s) **113**. In some instances, neural network host platform **110** may be made up of a plurality of different computing devices, which may be distributed within a single data center or a plurality of different data centers. In these instances, the one or more processor(s) **111**, one or more memory(s) **112**, and one or more communication interface(s) **113** included in neural network host platform **110** may be part of and/or otherwise associated with the different computing devices that form neural network host platform **110**.

(13) In one or more arrangements, processor(s) **111** may control operations of neural network host platform **110**. Memory(s) **112** may store instructions that, when executed by processor(s) **111**, cause neural network host platform **110** to perform one or more functions, as discussed below. Communication interface(s) **113** may include one or more wired and/or wireless network interfaces, and communication interface(s) **113** may connect neural network host platform **110** to one or more networks (e.g., network **190**) and/or enable neural network host platform **110** to exchange information and/or otherwise communicate with one or more devices connected to such

networks.

(14) In one or more arrangements, memory(s) **112** may store and/or otherwise provide a plurality of modules (which may, e.g., include instructions that may be executed by processor(s) **111** to cause neural network host platform **110** to perform various functions) and/or databases (which may, e.g., store data used by neural network host platform **110** in performing various functions). For example, memory(s) **112** may store and/or otherwise provide neural network host module **112a**, neural network host database **112b**, and a machine learning engine **112c**. In some instances, neural network host module **112a** may store instructions that cause neural network host platform **110** to apply machine learning for anomaly detection, and/or execute one or more other functions described herein. Additionally, neural network host database **112b** may store data that is used by neural network host platform **110** in applying machine learning for anomaly detection and/or in executing one or more other functions described herein. Furthermore, machine learning engine **112c** may store instructions and/or data that may cause and/or be used by neural network host platform **110** to identify anomalies in cybersecurity training modules and/or execute one or more other functions described herein.

(15) Cybersecurity module host system **120** may be and/or include one or more computing devices that may be configured to host one or more cybersecurity modules. For example, the cybersecurity module host system **120** may include one or more servers, server blades, or other devices configured for data storage. In some instances, in hosting the one or more cybersecurity modules, the cybersecurity module host system **120** may store one or more templates that may be used to generate various permutations of each cybersecurity module (e.g., based on language, device configuration, and/or other parameters).

(16) Enterprise user device **130** may be configured to be used by a first user (who may e.g., be an employee of an enterprise organization). In some instances, enterprise user device **130** may be configured to present one or more user interfaces associated with cybersecurity training modules, receive input corresponding to user interactions with the cybersecurity training modules, and/or otherwise facilitate participation in cybersecurity training modules.

(17) Administrator computing device **140** may be configured to be used by an administrative user (who may, e.g., be a network administrator and/or a cybersecurity analyst associated with an enterprise organization). Administrator computing device **140** may be configured to present one or more user interfaces associated with an operator dashboard, receive user input modifying training modules and/or templates for which an anomaly was detected, and/or otherwise facilitate monitoring and management of one or more systems and/or devices included in computing environment **100**.

(18) FIGS. 2A-2E depict an illustrative event sequence for applying machine learning for anomaly detection in accordance with one or more example embodiments. Referring to FIG. 2A, at step **201**, the neural network host platform **110** may train an auto-encoder to compute reconstruction error values. For example, the neural network host platform **110** may train the auto-encoder to compute values indicating a discrepancy between actual permutations of training modules (e.g., screenshots of permutations of the training modules, which may be generated by the neural network host platform **110**) and anticipated permutations of the training modules (e.g., screenshots indicating what the permutations of the training modules should look like). To do so, the neural network host platform **110** may train the auto-encoder using one or more unsupervised learning techniques. In some instances, the neural network host platform **110** may use screenshots of cybersecurity module permutations configured for different computing devices (e.g., laptop computers, desktop computers, mobile devices, and/or other devices), different languages, different web browsers, different resolutions, and/or other configurable parameters/features to train the auto-encoder. In some instances, the neural network host platform **110** may use screenshots of entire cybersecurity module permutations and/or isolated visual components from cybersecurity module permutations to train the auto-encoder.

(19) At step **202**, the administrator computing device **140** may receive a selection input (e.g., from an administrative user and/or cybersecurity analyst operating the administrator computing device **140**). In some instances, in receiving the selection input, the administrator computing device **140** may receive an input selecting one or more cybersecurity modules for testing (e.g., to be tested for anomalies in various pages of various permutations of the selected one or more cybersecurity modules). In some instances, the administrator computing device **140** may receive the selection input through one or more portals or other graphical user interfaces.

(20) At step **203**, administrator computing device **140** may send, share, or otherwise provide a testing request to the neural network host platform **110**. For example, based on or in response to the selection input received at step **202**, the administrator computing device **140** may send a request to the neural network host platform **110** to test (e.g., for uniformity) the selected one or more cybersecurity modules. In some instances, in sending the testing request, the neural network host platform **110** may send an identifier that may be used to identify the selected one or more training modules. Additionally or alternatively, the administrator computing device **140** may send one or more commands directing the neural network host platform to test the selected one or more cybersecurity modules.

(21) At step **204**, the neural network host platform **110** may receive or otherwise access the testing request from the administrator computing device **140** sent at step **203**. In some instances, in receiving the testing request, the neural network host platform **110** may receive an identifier that may be used to identify the selected one or more training modules. Additionally or alternatively, the neural network host platform **110** may receive the one or more commands directing the neural network host platform **110** to test the selected one or more training modules.

(22) At step **205**, the neural network host platform **110** may generate a request for cybersecurity module templates and/or template information corresponding to the selected one or more cybersecurity modules (which were identified in the testing request received at step **204**). For example, the neural network host platform **110** may generate a request to the cybersecurity module host system **120** to provide templates corresponding to the one or more selected cybersecurity modules. In some instances, the neural network host platform **110** may include the identifier of the selected one or more cybersecurity requests in the request. Additionally or alternatively, the neural network host platform **110** may send one or more commands directing the cybersecurity module host system **120** to provide the cybersecurity module templates.

(23) Referring to FIG. 2B, at step **206**, the neural network host platform **110** may send, share, or otherwise provide the request for cybersecurity module templates and/or template information (generated at step **205**) to the cybersecurity module host system **120**. Additionally or alternatively, the neural network host platform **110** may send the one or more commands directing the cybersecurity module host system **120** to provide the cybersecurity module templates.

(24) At step **207**, the cybersecurity module host system **120** may receive or otherwise access the request for the cybersecurity module templates and/or template information, sent at step **206**. Additionally or alternatively, the cybersecurity module host system **120** may receive the one or more commands directing the cybersecurity module host system **120** to provide the cybersecurity module templates.

(25) At step **208**, the cybersecurity module host system **120** may send, share, or otherwise provide the cybersecurity module templates and/or template information based on or in response to the request for the cybersecurity module templates received at step **207**. For example, the cybersecurity module host system **120** may identify cybersecurity module templates, based on the identification information included in the request for the cybersecurity module templates, that may be used by the neural network host platform **110** to produce different permutations of the selected one or more cybersecurity training modules (e.g., information that may define the selected one or more cybersecurity modules).

(26) At step **209**, the neural network host platform **110** may receive or otherwise access the

cybersecurity module templates and/or template information sent at step **208**. For example, the neural network host platform **110** may receive the cybersecurity module templates corresponding to the selected one or more cybersecurity modules (e.g., information that may define the selected one or more cybersecurity modules).

(27) At step **210**, the neural network host platform **110** may render permutations of the selected one or more cybersecurity modules using the cybersecurity module templates and/or template information received at step **209**. For example, the neural network host platform **110** may render permutations corresponding to versions of the selected one or more cybersecurity modules in various languages, configured for display at various devices (e.g., configured for particular devices, operating systems, or otherwise based on device parameters), configured for display in various browsers, and/or otherwise configured based on additional parameters. In doing so, the neural network host platform **110** may render each page of each of the selected one or more cybersecurity modules for each of the permutations. For example, the neural network host platform **110** may render, for each permutation, a plurality of pages that may be used by an enterprise user (e.g., user of enterprise user device **130**) to progress through a cybersecurity module (e.g., one of the selected one or more cybersecurity modules). In some instances, the neural network host platform **110** may store each rendered permutation as well as distinct elements and/or other portions of each rendering. In some instances, the neural network host platform **110** may receive an input indicating visual elements, attributes, and/or other features of interest (e.g., from a system administrator) that should be distinctly stored. In these instances, the neural network host platform **110** may access information indicating these visual elements, attributes, and/or other features of interest, which may be manually coded and/or otherwise embedded into the selected one or more cybersecurity training modules.

(28) At step **211**, the neural network host platform **110** may capture screenshots for each of the permutations rendered at step **210**. For example, the neural network host platform **110** may capture a screenshot for each page within each cybersecurity module for which permutations were rendered at step **210**. As a particular example, the neural network host platform **110** may capture screenshots for each page of a “Security Essentials” training module that are each configured in a plurality of different languages, browser configurations, and/or other formats. In some instances, in capturing the screenshots, the neural network host platform **110** may capture a screenshot of a permutation and may crop out various reusable components. In these instances, the screenshots referred to in the proceeding steps may refer to these cropped components. In doing so, the neural network host platform **110** may model an entire training module/permutation and/or individual components that make up the training modules/permutations.

(29) Referring to FIG. 2C, at step **212**, the neural network host platform **110** may pre-process the screenshots captured at step **211**. For example, the neural network host platform **110** may redact and/or otherwise mask text, images, and/or other components from each screenshot. In doing so, the neural network host platform **110** may improve identification of errors in visibility, size, positioning of specific elements, and/or other formatting errors. In some instances, the neural network host platform **110** may redact each type of visual component with a unique color (e.g., first color for text and a second color for images). Additionally or alternatively, the neural network host platform **110** may redact each visual component with a unique color, which may preserve parent/child element fields by apply a unique color to conceal both the parent element and the child element. In instances where the neural network host platform **110** redacts text from the screenshots, the neural network host platform **110** may improve effectiveness of layout uniformity testing across the various screenshots. In instances where the neural network host platform **110** redacts images from the screenshots, the neural network host platform may improve uniformity analysis of visibility, size, position, and/or other image properties for the screenshots. In instances where the neural network host platform **110** redacts other components, the neural network host platform **110** may improve layout uniformity analysis.

(30) At step **213**, the neural network host platform **110** may input the pre-processed screenshots into the auto-encoder trained at step **201**. In some instances, the auto-encoder may be hosted by the neural network host platform **110** or another computing system with which the neural network host platform **110** is configured to communicate. In some instances, inputting the pre-processed screenshots into the auto-encoder may cause the neural network host platform **110** to output the reconstruction error values as described below at step **214**.

(31) At step **214**, the neural network host platform **110** may use the auto-encoder to identify reconstruction error values for each screenshot. For example, the neural network host platform **110** may use the auto-encoder to compare the pre-processed screenshots to anticipated screenshots for a corresponding cybersecurity module (which may e.g., have been used to train the auto-encoder at step **201**). Additionally or alternatively, the neural network host platform **110** may use the auto-encoder to compare the pre-processed screenshots to each other (e.g., identify a format discrepancy between a screenshot of a first permutation of a particular cybersecurity module and the corresponding screenshots for the remaining permutations of the particular cybersecurity module). In identifying these reconstruction error values, the neural network host platform **110** may use the auto-encoder to output error values indicating a degree of discrepancy between a captured screenshot of a cybersecurity module (e.g., which in some instances may be a screenshot of an outlier permutation of the cybersecurity module) and a model generated by the auto-encoder training process. In these instances, the neural network host platform **110** may identify lower reconstruction error values where the degree of discrepancy is lower and higher reconstruction error values where the degree of discrepancy is higher.

(32) In some instances, in identifying the reconstruction error values, the neural network host platform **110** may assign a numeric value for each identified discrepancy. In some instances, the neural network host platform **110** may assign a consistent numeric value for each identified discrepancy. In other instances, the neural network host platform **110** may assign different numeric values for different discrepancies (e.g., first numeric value for misplaced element and second numeric value for incorrect element).

(33) In some instances, the neural network host platform **110** may compare the reconstruction error values to a predetermined reconstruction error threshold. For reconstruction error values, identified by the neural network host platform **110** to be greater than the predetermined reconstruction error threshold, the neural network host platform **110** may proceed to step **215** to identify outlier permutations. For reconstruction error values, identified by the neural network host platform **110** to be less than or equal to the predetermined reconstruction error threshold, the neural network host platform **110** may proceed to step **224**. In doing so, the neural network host platform **110** may effectively filter permutations for which the outlier detection algorithm does not need to be applied at step **215**, which may conserve computing resources.

(34) At step **215**, the neural network host platform **110** may identify, based on the reconstruction error values output at step **214**, one or more outlier permutations. In some instances, the neural network host platform **110** may identify the one or more outlier permutations by executing an outlier detection algorithm using the reconstruction error values. For example, the neural network host platform **110** may compute, based on the reconstruction error values corresponding to different permutations for a particular page of a particular cybersecurity module, a standard deviation for the particular page of the particular cybersecurity module. For example, the neural network host platform **110** may apply the following outlier detection algorithm

$$(35) \quad \sigma = \sqrt{\frac{(x_i - \mu)^2}{N}},$$

where σ is the standard deviation, N is the number of permutations of the particular page of the particular cybersecurity module, $x_{sub.i}$ is the reconstruction error value for each of the permutations of the particular page of the particular cybersecurity module, and μ is the mean reconstruction error value for the permutations of the particular page of the particular cybersecurity

module.

(36) In some instances, the neural network host platform **110** may compute the standard deviation on a page by page basis (e.g., compare corresponding pages of different permutations of the same cybersecurity module). Additionally or alternatively, the neural network host platform **110** may compute the standard deviation on an element by element basis (e.g., compare corresponding elements of corresponding pages of different permutations of the same cybersecurity module). In these instances, the neural network host platform **110** may use these element by element deviations to compute an overall standard deviation for the page (e.g., by averaging or otherwise combining the element by element standard deviations).

(37) After computing this standard deviation, the neural network host platform **110** may compare the standard deviation to a predetermined outlier identification threshold. In this example, if the neural network host platform **110** determines that the standard deviation exceeds the predetermined outlier identification threshold, the neural network host platform **110** may determine that the particular page of the particular cybersecurity module includes at least one outlier permutation. For example, the neural network host platform **110** may determine that the corresponding cybersecurity module contains one or more anomalies, and should be updated to correct them. In these instances, the neural network host platform **110** may proceed to step **216**. If the neural network host platform **110** determines that the standard deviation is less than or equal to the predetermined outlier identification threshold, the neural network host platform **110** may determine that the corresponding cybersecurity module does not include an outlier permutation. For example, the neural network host platform **110** may determine that the corresponding cybersecurity module does not need to be updated to correct anomalies (e.g., the anomalies are minimal and/or non-existent). In these instances, the event sequence may proceed to step **224**. In some instances, in comparing the standard deviation to the predetermined outlier identification threshold, the neural network host platform **110** to a dynamic/variable threshold (e.g., the predetermined outlier identification threshold might not be a fixed value). In these instances, the neural network host platform **110** may use the predetermined outlier identification threshold to provide a “yes” or “no” determination for whether or not an outlier permutation is identified.

(38) In some instances, in addition or as an alternative to comparing the standard deviation to the predetermined outlier identification threshold, the neural network host platform **110** may select a predetermined number of cybersecurity modules with the highest reconstruction error values (e.g., the cybersecurity modules with the ten (or some other predetermined number) highest reconstruction error values).

(39) Referring to FIG. 2D, at step **216**, the neural network host platform **110** may generate an outlier interface indicating any cybersecurity modules, identified at step **215**, that should be updated or otherwise modified. In some instances, in generating the outlier interface, the neural network host platform **110** may generate an interface that may allow a user of the administrator computing device **140** to edit/correct outlier permutations identified at step **215**.

(40) At step **217**, the neural network host platform **110** may send, share, or otherwise provide the outlier interface, generated at step **217**, to the administrator computing device **140**. At step **218**, the administrator computing device **140** may receive or otherwise access the outlier interface, sent at step **217**. At step **219**, the administrator computing device **140** may display the outlier interface received at step **218**. In some instances, sending the outlier interface at step **217** may cause the administrator computing device **140** to display the outlier interface at step **219**. In some instances, in displaying the outlier interface, the administrator computing device **140** may display a graphical user interface similar to graphical user interface **300**, which is shown in FIG. 3. For example, the administrator computing device **140** may display a graphical user interface that indicates, using for example a numeric identifier, alphabetic identifier, and/or alpha-numeric identifier, cybersecurity modules with detected anomalies, a specific page of the cybersecurity modules with the detected anomalies, and/or a specific permutation of the cybersecurity modules with the detected anomalies.

Additionally or alternatively, the administrator computing device may display, on the graphical user interface, a summary of the anomaly, a summary of corrections to be made, a status of remedying the anomaly, an individual responsible for addressing the anomaly, and/or other information related to detected anomalies.

(41) At step **220**, the administrator computing device **140** may receive a template modification input (e.g., from a user of the administrator computing device **140**). For example, the administrator computing device **140** may receive an input modifying one or more templates, corresponding to the cybersecurity modules for which anomalies were detected, that may be used to generate these cybersecurity modules and/or permutations of the cybersecurity modules. In some instances, in receiving the template modification input, the administrator computing device **140** may receive input recoding the cybersecurity modules. In these examples, the administrator computing device **140** may receive an input that may address the detected anomalies, and may result in uniformity between different permutations of each cybersecurity module for which anomalies were detected. In some instances, the administrator computing device **140** may receive an input indicating whether or not an identified outlier permutation was correctly identified as an outlier permutation. For example, the administrator computing device **140** may receive selection of a thumbs up or thumbs down icon, which may indicate that the identified outlier permutation was or was not correctly identified as an outlier permutation respectively. In these instances, the administrator computing device **140** may send feedback information to the neural network host platform **110**, and the neural network host platform **110** may dynamically tune the auto-encoder based on the feedback information.

(42) Referring to FIG. 2E, at step **221**, the administrator computing device **140** may send, share, or otherwise provide a template modification information, based on the template modification input received at step **220**, to the cybersecurity module host system **120**. For example, the administrator computing device **140** may send information to the cybersecurity module host system **120** that may cause the cybersecurity module host system **120** to correct identified anomalies between corresponding pages of various renderings of stored cybersecurity modules.

(43) At step **222**, the cybersecurity module host system **120** may receive or otherwise access the template modification information sent at step **221**. For example, the cybersecurity module host system **120** may receive information that may cause the cybersecurity module host system **120** to correct identified anomalies between corresponding pages of various renderings of stored cybersecurity modules.

(44) At step **223**, the cybersecurity module host system **120** may update cybersecurity module templates based on the template modification information received at step **222**. For example, the cybersecurity module host system **120** may update stored cybersecurity module templates to address anomalies in identified outlier permutations (e.g., at steps **214-215**).

(45) At step **224**, the cybersecurity module host system **120** may send, share, or otherwise provide a cybersecurity module to the enterprise user device **130**. For example, the cybersecurity module host system **120** may send one of the cybersecurity modules for which a template was updated at step **223**.

(46) At step **225**, the enterprise user device **130** may receive or otherwise access the cybersecurity module sent at step **224**. For example, the enterprise user device **130** may receive a cybersecurity module for which a template was updated at step **223**.

(47) At step **226**, the enterprise user device **130** may display the cybersecurity module received at step **226** (which may e.g., have been updated at step **223** to correct existing anomalies). For example, sending the cybersecurity module at step **224** may cause the enterprise user device **130** to display the cybersecurity module at step **226**. In some instances, in displaying the cybersecurity module, the enterprise user device **130** may display a graphical user interface similar to graphical user interface **400**, which is shown in FIG. 4. For example, the enterprise user device **130** may display a series of pages that a user of the enterprise user device **130** (e.g., an employee of an

enterprise organization) may progress through to gain cybersecurity training (e.g., a security essentials module for executives of the enterprise organization, or other modules). In some instances, in displaying the cybersecurity module, the enterprise user device **130** may display a graphical user interface that prompts for feedback indicating whether or not the cybersecurity module includes any visible anomalies and/or other feedback. In these instances, the enterprise user device **130** may send any received feedback to the neural network host platform **110**, and the neural network host platform **110** may use this feedback to dynamically tune the auto-encoder. Additionally or alternatively, the neural network host platform **110** may use the feedback to filter future results. For example, the neural network host platform **110** may effectively “mute” outliers identified as a non-issue by storing a hash of the corresponding screenshot and subsequently checking future anomaly hashes to identify whether or not the anomaly should be suppressed when reporting identified issues.

(48) The steps described in the illustrative event sequence herein may be performed in any alternative sequence or order without departing from the scope of the disclosure. Furthermore, although the above described systems, event sequence, and methods are described primarily in the context of cybersecurity training, this is solely for illustrative purposes, and the systems, event sequence, and methods described herein may be applied in other contexts (e.g., maintaining uniformity/anomaly detection for any type of module(s), web page(s), user interface(s), and/or other displayable content) without departing from the scope of the disclosure. In some instances, these other contexts may include internal enterprise activities and/or external activities (e.g., outside of a particular enterprise organization). Furthermore, although the above described systems, event sequence, and methods are described in the context of internal cybersecurity training, the systems, event sequence, and methods described herein may be used to provide a software as a service (SaaS) visual testing service for localized pages corresponding to various other products without departing from the scope of the disclosure.

(49) FIG. 5 depicts an illustrative method for applying machine learning in anomaly detection in accordance with one or more example embodiments. Referring to FIG. 5, at step **505**, a computing platform having at least one processor, a communication interface, and memory may train an auto-encoder for anomaly detection in cybersecurity modules. At step **510**, the computing platform may receive a request to test one or more cybersecurity modules. At step **515**, the computing platform may generate and send a request for cybersecurity templates corresponding to the one or more cybersecurity modules. At step **520**, the computing platform may receive the cybersecurity training module templates. At step **525**, the computing platform may render permutations of the cybersecurity training module templates. At step **530**, the computing platform may capture screenshots for each page of each rendered permutation of the cybersecurity training module templates. At step **535**, the computing platform may pre-process the screenshots captured at step **530**. At step **540**, the computing platform may input the screenshots into the auto-encoder. At step **545**, the computing platform may use the auto-encoder to output reconstruction error values. At step **550**, the computing platform may determine whether or not any of the reconstruction error values correspond to an outlier permutation. If the computing platform does not identify one or more outlier permutations, the method may end. If the computing platform does identify one or more outlier permutations, the computing platform may proceed to step **555**. At step **555**, the computing platform may generate and send an outlier interface to an administrator computing device.

(50) It should be understood that the analysis processes, method steps, and/or methods described herein may be performed in different orders and/or in alternative arrangements from those illustrated herein, without departing from the scope of this disclosure. Additionally or alternatively, one or more of the analysis processes, method steps, and/or methods described herein may be optional and/or omitted in some arrangements, without departing from the scope of this disclosure.

(51) One or more aspects of the disclosure may be embodied in computer-usable data or computer-

executable instructions, such as in one or more program modules, executed by one or more computers or other devices to perform the operations described herein. Program modules may include routines, programs, objects, components, data structures, and the like that perform particular tasks or implement particular abstract data types when executed by one or more processors in a computer or other data processing device. The computer-executable instructions may be stored as computer-readable instructions on a computer-readable medium such as a hard disk, optical disk, removable storage media, solid-state memory, RAM, and the like. The functionality of the program modules may be combined or distributed as desired in various embodiments. In addition, the functionality may be embodied in whole or in part in firmware or hardware equivalents, such as integrated circuits, application-specific integrated circuits (ASICs), field programmable gate arrays (FPGA), and the like. Particular data structures may be used to more effectively implement one or more aspects of the disclosure, and such data structures are contemplated to be within the scope of computer executable instructions and computer-usable data described herein.

(52) One or more aspects described herein may be embodied as a method, an apparatus, or as one or more computer-readable media storing computer-executable instructions. Accordingly, those aspects may take the form of an entirely hardware embodiment, an entirely software embodiment, an entirely firmware embodiment, or an embodiment combining software, hardware, and firmware aspects in any combination. In addition, various signals representing data or events as described herein may be transferred between a source and a destination in the form of light or electromagnetic waves traveling through signal-conducting media such as metal wires, optical fibers, or wireless transmission media (e.g., air or space). The one or more computer-readable media may be and/or include one or more non-transitory computer-readable media.

(53) As described herein, the various methods and acts may be operative across one or more computing servers and one or more networks. The functionality may be distributed in any manner, or may be located in a single computing device (e.g., a server, a client computer, and the like). For example, in alternative embodiments, one or more of the computing platforms discussed above may be combined into a single computing platform, and the various functions of each computing platform may be performed by the single computing platform. In such arrangements, any and/or all of the above-discussed communications between computing platforms may correspond to data being accessed, moved, modified, updated, and/or otherwise used by the single computing platform. Additionally or alternatively, one or more of the computing platforms discussed above may be implemented in one or more virtual machines that are provided by one or more physical computing devices. In such arrangements, the various functions of each computing platform may be performed by the one or more virtual machines, and any and/or all of the above-discussed communications between computing platforms may correspond to data being accessed, moved, modified, updated, and/or otherwise used by the one or more virtual machines.

(54) Aspects of the disclosure have been described in terms of illustrative embodiments thereof. Numerous other embodiments, modifications, and variations within the scope and spirit of the appended claims will occur to persons of ordinary skill in the art from a review of this disclosure. For example, one or more of the steps depicted in the illustrative figures may be performed in other than the recited order, and one or more depicted steps may be optional in accordance with aspects of the disclosure.

Claims

1. A computing platform, comprising: at least one processor; a communication interface communicatively coupled to the at least one processor; and memory storing computer-readable instructions that, when executed by the at least one processor, cause the computing platform to: receive information defining a training module; capture a plurality of screenshots corresponding to

different permutations of the training module; input, into an auto-encoder, the screenshots corresponding to the different permutations of the training module, wherein inputting the screenshots corresponding to the different permutations of the training module causes the auto-encoder to compare the screenshots to anticipated screenshots for the training module to output a reconstruction error value; identify, based on the reconstruction error value, an outlier permutation of the training module; generate a user interface comprising information identifying the outlier permutation of the training module; and send the user interface to at least one user device.

2. The computing platform of claim 1, wherein the auto-encoder is trained on a set of images corresponding to a plurality of training modules including the training module.

3. The computing platform of claim 2, wherein the set of images includes images configured for one or more of: a desktop computer, a laptop computer, or a mobile device.

4. The computing platform of claim 1, wherein the user interface includes controls allowing a user of the at least one user device to edit the outlier permutation of the training module.

5. The computing platform of claim 1, wherein the different permutations of the training module correspond to one or more of: different languages, different browsers, or different resolutions.

6. The computing platform of claim 1, wherein the memory stores additional computer-readable instructions that, when executed by the at least one processor, cause the computing platform to: render, using the information defining the training module, the different permutations of the training module.

7. The computing platform of claim 1, wherein sending the user interface to the at least one user device causes the at least one user device to display the user interface.

8. The computing platform of claim 1, wherein the user interface includes one or more of: an alphanumeric identifier for the outlier permutation of the training module, a link to access the outlier permutation of the training module, a summary of corrections to be made to the outlier permutation of the training module, an individual assigned to perform the corrections, or a status of the corrections.

9. The computing platform of claim 1, wherein the reconstruction error value indicates a degree to which the outlier permutation of the training module differs from an anticipated permutation of the training module.

10. The computing platform of claim 9, wherein identifying the outlier permutation of the training module comprises: identifying, using an outlier detection algorithm and the reconstruction error value, that the degree to which the outlier permutation of the training module differs from the anticipated permutation of the training module exceeds a predetermined anomaly identification threshold; and based on the identification that the degree to which the outlier permutation of the training module differs from the anticipated permutation of the training module exceeds the predetermined anomaly identification threshold, identifying the outlier permutation of the training module.

11. The computing platform of claim 1, wherein the memory stores additional computer-readable instructions that, when executed by the at least one processor, cause the computing platform to: receive, after sending the user interface to the at least one user device, user input indicating whether or not the outlier permutation of the training module was correctly identified as an outlier; and dynamically tune, based on the user input indicating whether or not the outlier permutation of the training module was correctly identified as an outlier, the auto-encoder.

12. A method comprising: at a computing platform comprising at least one processor, a communication interface, and memory: receiving, by the at least one processor, information defining a training module; capturing, by the at least one processor, a plurality of screenshots corresponding to different permutations of the training module; inputting, by the at least one processor and into an auto-encoder, the screenshots corresponding to the different permutations of the training module, wherein inputting the screenshots corresponding to the different permutations of the training module causes the auto-encoder to compare the screenshots to anticipated

screenshots for the training module to output a reconstruction error value; identifying, by the at least one processor and based on the reconstruction error value, an outlier permutation of the training module; generating, by the at least one processor, a user interface comprising information identifying the outlier permutation of the training module; and sending, by the at least one processor, the user interface to at least one user device.

13. The method of claim 12, wherein the auto-encoder is trained on a set of images corresponding to a plurality of training modules including the training module.

14. The method of claim 13, wherein the set of images includes images configured for one or more of: a desktop computer, a laptop computer, or a mobile device.

15. The method of claim 12, wherein the user interface includes controls allowing a user of the at least one user device to edit the outlier permutation of the training module.

16. The method of claim 12, wherein the different permutations of the training module correspond to one or more of: different languages, different browsers, or different resolutions.

17. The method of claim 12, further comprising: rendering, by the at least one processor and using the information defining the training module, the different permutations of the training module.

18. The method of claim 12, wherein sending the user interface to the at least one user device causes the at least one user device to display the user interface.

19. The method of claim 12, wherein the user interface includes one or more of: an alpha-numeric identifier for the outlier permutation of the training module, a link to access the outlier permutation of the training module, a summary of corrections to be made to the outlier permutation of the training module, an individual assigned to perform the corrections, or a status of the corrections.

20. One or more non-transitory computer-readable media storing instructions that, when executed by a computing platform comprising at least one processor, a communication interface, and memory, cause the computing platform to: receive information defining a training module; capture a plurality of screenshots corresponding to different permutations of the training module; input, into an auto-encoder, the screenshots corresponding to the different permutations of the training module, wherein inputting the screenshots corresponding to the different permutations of the training module causes the auto-encoder to compare the screenshots to anticipated screenshots for the training module to output a reconstruction error value; identify, based on the reconstruction error value, an outlier permutation of the training module; generate a user interface comprising information identifying the outlier permutation of the training module; and send the user interface to at least one user device.
