

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12387475
B2
August 12, 2025
Casas; Sergio et al.

Systems and methods for latent distribution modeling for scene-consistent motion forecasting

Abstract

A computer-implemented method for determining scene-consistent motion forecasts from sensor data can include obtaining scene data including one or more actor features. The computer-implemented method can include providing the scene data to a latent prior model, the latent prior model configured to generate scene latent data in response to receipt of scene data, the scene latent data including one or more latent variables. The computer-implemented method can include obtaining the scene latent data from the latent prior model. The computer-implemented method can include sampling latent sample data from the scene latent data. The computer-implemented method can include providing the latent sample data to a decoder model, the decoder model configured to decode the latent sample data into a motion forecast including one or more predicted trajectories of the one or more actor features. The computer-implemented method can include receiving the motion forecast including one or more predicted trajectories of the one or more actor features from the decoder model.

Inventors: Casas; Sergio (Toronto, CA), Gulino; Cole Christian (Pittsburgh, PA), Suo; Shun Da (Toronto, CA), Luo; Katie Z. (Sunnyvale, CA), Liao; Renjie (Toronto, CA), Urtasun; Raquel (Toronto, CA)

Applicant: Aurora Operations, Inc. (Pittsburgh, PA)

Family ID: 77554739

Assignee: AURORA OPERATIONS, INC. (Pittsburgh, PA)

Appl. No.: 18/519976

Filed: November 27, 2023

Prior Publication Data

Document Identifier
US 20240096083 A1

Publication Date
Mar. 21, 2024

Related U.S. Application Data

continuation parent-doc US 17150995 20210115 US 11842530 child-doc US 18519976
us-provisional-application US 63119981 20201201
us-provisional-application US 62985862 20200305

Publication Classification

Int. Cl.: **G06V10/82** (20220101); **B60W60/00** (20200101); **G05D1/00** (20240101); **G06F18/2137** (20230101); **G06V20/58** (20220101); **G06V30/19** (20220101); **G06V30/24** (20220101); **G08G1/16** (20060101)

U.S. Cl.:

CPC **G06V10/82** (20220101); **B60W60/0027** (20200201); **G05D1/0212** (20130101); **G06F18/2137** (20230101); **G06V20/58** (20220101); **G06V30/19173** (20220101); **G08G1/166** (20130101); G06V30/2504 (20220101)

Field of Classification Search

CPC: G06V (10/82); G06V (20/58); G06V (30/2504); G06V (30/19173); G08G (1/166); G05D (1/0212); G06F (18/2137); B60W (60/0027)

USPC: 701/25

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11618162	12/2022	Cristache	700/245	B25J 9/1697
2019/0050648	12/2018	Stojanovic	N/A	G06V 20/13
2019/0135300	12/2018	Gonzalez Aguirre	N/A	G06N 20/00
2019/0339770	12/2018	Kurlethimar	N/A	G02B 27/017
2020/0234066	12/2019	Lee	N/A	G05D 1/0251
2020/0293041	12/2019	Palanisamy	N/A	G06N 3/08
2020/0327674	12/2019	Yang et al.	N/A	N/A
2021/0276595	12/2020	Casas et al.	N/A	N/A
2022/0075383	12/2021	Morad	N/A	G06N 3/084
2022/0417428	12/2021	Jeon	N/A	H04N 23/69

OTHER PUBLICATIONS

Alahi et al, "Social LSTM: Human Trajectory Prediction in Crowded Spaces", Conference on Computer Vision and Pattern Recognition, Jun. 26-Jul. 1, 2016, Las Vegas, Nevada, United States, 11 pages. cited by applicant

Bengio et al, "Scheduled Sampling for Sequence Prediction with Recurrent Neural Networks", arXiv:1506.03099v3, Sep. 23, 2015, 9 pages. cited by applicant

Caesar et al, "NuScenes: A Multimodal Dataset for Autonomous Driving", arXiv:1903.11027v5, May 5, 2020, 16 pages. cited by applicant

Casas et al, "IntentNet: Learning to Predict Intention from Raw Sensor Data",

arXiv:2101.07907v1, Jan. 20, 2021, 10 pages. cited by applicant

Casas et al, "SpAGNN: Spatially-Aware Graph Neural Networks for Relational Behavior Forecasting from Sensor Data", arXiv:1910.08233v1, Oct. 18, 2019, 11 pages. cited by applicant

Chai et al, "Multipath: Multiple Probabilistic Anchor Trajectory Hypotheses for Behavior Prediction", arXiv:1910.05449v1, Oct. 12, 2019, 14 pages. cited by applicant

Cui et al, "Multimodal Trajectory Predictions for Autonomous Driving Using Deep Convolutional Networks", arXiv:1809.10732v2, Mar. 1, 2019, 7 pages. cited by applicant

Djuric et al, "Motion Prediction of Traffic Actors for Autonomous Driving Using Deep Convolutional Networks", arXiv:1808.05819v1, Aug. 17, 2018, 10 pages. cited by applicant

Higgins et al, " β -VAE: Learning Basic Visual Concepts with a Constrained Variational Framework", International Conference on Learning Representations, Apr. 24-26, 2017, Toulon, France, 22 pages. cited by applicant

Hong et al, "Rules of the Road: Predicting Driving Behavior with a Convolutional Model of Semantic Interactions", arXiv:1906.08945v1, Jun. 21, 2019, 9 pages. cited by applicant

Hubmann et al, "Automated Driving in Uncertain Environments: Planning with Interaction and Uncertain Maneuver Prediction" Transactions on Intelligent Vehicles vol. 3, No. 1, Mar. 2018, pp. 5-17. cited by applicant

Huszar, "How (Not) to Train your Generative Model: Scheduled Sampling, Likelihood, Adversary?", arXiv:1511.05101v1, Nov. 16, 2015, 9 pages. cited by applicant

Jain et al, "Discrete Residual Flow for Probabilistic Pedestrian Behavior Prediction", arXiv:1910.08041v1, Oct. 17, 2019, 13 pages. cited by applicant

Jang et al, "Categorical Reparameterization with Gumbel-Softmax", arXiv:1611.01144v5, Aug. 5, 2017, 13 pages. cited by applicant

Kingma et al, "Auto-Encoding Variational Bayes", arXiv:1312.6114v10, May 1, 2014, 14 pages. cited by applicant

Klingelschmitt et al, "Managing the Complexity of Inner-City Scenes: An Efficient Situation Hypotheses Selection Scheme", Intelligent Vehicles Symposium (IV), 2015, 9 pages. cited by applicant

Lamb et al, "Professor Forcing: A New Algorithm for Training Recurrent Networks", arXiv:1610.09038v1, Oct. 27, 2016, 11 pages. cited by applicant

Lang et al, PointPillars: Fast Encoders for Object Detection from Point Clouds, arXiv:1812.05784v2, May 7, 2019, 9 pages. cited by applicant

Lee et al, "Desire: Distant Future Prediction in Dynamic Scenes with Interacting Agents", arXiv:1704.04394v1, Apr. 14, 2017, 10 pages. cited by applicant

Liang et al, "Multi-Task Multi-Sensor Fusion for 3D Object Detection", arXiv:2012.12397v1, Dec. 22, 2020, 11 pages. cited by applicant

Luo et al, "Fast and Furious: Real Time End-to-End 3D Detection, Tracking and Motion Forecasting with A Single Convolutional Net", arXiv:2012.12395v1, Dec. 22, 2020, 9 pages. cited by applicant

Ma et al, "Arbitrary Oriented Scene Text Detection via Rotation Proposals", arXiv:1703.01086v3, Mar. 15, 2018, 11 pages. cited by applicant

Okamoto et al, "DDT: Deep Driving Tree for Proactive Planning in Interactive Scenarios" In: 2018 21st International Conference on Intelligent Transportation Systems, Nov. 4-7, 2018, Maui, Hawaii, United States, pp. 656-661. cited by applicant

Osa, et al, "An Algorithmic Perspective on Imitation Learning", arXiv:1811.06711v1, Nov. 16, 2018, 188 pages. cited by applicant

Qi et al, "PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation", arXiv:1612.00593v2, Apr. 10, 2017, 19 pages. cited by applicant

Ratliff et al, "Maximum Margin Planning", International Conference on Machine Learning, Jun. 25-29, 2006, Washington, United States, 8 pages. cited by applicant

Rhinehart et al, "PRECOG: PREdiction Conditioned on Goals in Visual Multi-Agent Settings", arXiv:1905.01296v3, Sep. 30, 2019, 24 pages. cited by applicant

Rhinehart et al, "R2P2: A Reparameterized Pushforward Policy for Diverse, Precise Generative Path Forecasting", European Conference on Computer Vision, Sep. 8-14, 2018, Munich, Germany, 17 pages. cited by applicant

Ross et al, "A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning", arXiv:1011.0686v3, Mar. 16, 2011, 9 pages. cited by applicant

Sadat et al, "Jointly Learnable Behavior and Trajectory Planning for Self-Driving Vehicles", arXiv:1910.04586v1, Oct. 10, 2019, 8 pages. cited by applicant

Schlichtkrull et al, "Modeling Relational Data with Graph Convolutional Networks", arXiv:1703.06103v4, Oct. 26, 2017, 9 pages. cited by applicant

Sohn et al, "Learning Structured Output Representation Using Deep Conditional Generative Models", Conference on Neural Information Processing Systems, Dec. 7-12, 2015, Montreal, Canada, 9 pages. cited by applicant

Tang et al, "Multiple Futures Prediction", arXiv:1911.00997v2, Dec. 6, 2019, 17 pages. cited by applicant

Yang et al, "PIXOR: Real-Time 3D Object Detection from Point Clouds", arXiv:1902.06326v3, Mar. 2, 2019, 10 pages. cited by applicant

Zeng et al, "End-to-End Interpretable Neural Motion Planner", arXiv:2101.06679v1, Jan. 17, 2021, 10 pages. cited by applicant

Zhou et al, "End-to-End Multi-View Fusion for 3D Object Detection in Lidar Point Clouds", arXiv:1910.06528v2, Oct. 23, 2019, 10 pages. cited by applicant

Primary Examiner: Soofi; Yazan A

Attorney, Agent or Firm: Dority & Manning, P.A.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application is based on and claims benefit of U.S. patent application Ser. No. 17/150,995 having a filing date of Jan. 15, 2021, U.S. Provisional Patent Application No. 63/119,981 having a filing date of Dec. 1, 2020, and U.S. Provisional Patent Application No. 62/985,862 having a filing date of Mar. 5, 2020, each both of which are incorporated by reference herein.

FIELD

(1) The present disclosure relates generally to autonomous vehicles. More particularly, the present disclosure relates to systems and methods for latent distribution modeling for scene-consistent motion forecasting.

BACKGROUND

(2) An autonomous vehicle is a vehicle that is capable of sensing its environment and navigating without human input. In particular, an autonomous vehicle can observe its surrounding environment using a variety of sensors and can attempt to comprehend the environment by performing various processing techniques on data collected by the sensors. Given knowledge of its surrounding environment, the autonomous vehicle can identify an appropriate motion path for navigating through such surrounding environment.

SUMMARY

(3) Aspects and advantages of embodiments of the present disclosure will be set forth in part in the following description, or can be learned from the description, or can be learned through practice of

the embodiments.

(4) One example aspect of the present disclosure is directed to a computer-implemented method for determining scene-consistent motion forecasts from sensor data. The computer-implemented method can include obtaining, by a computing system including one or more computing devices, scene data including one or more actor features. The computer-implemented method can include providing, by the computing system, the scene data to a latent prior model, the latent prior model configured to generate scene latent data in response to receipt of scene data, the scene latent data including one or more latent variables. The computer-implemented method can include obtaining, by the computing system, the scene latent data from the latent prior model. The computer-implemented method can include sampling, by the computing system, latent sample data from the scene latent data. The computer-implemented method can include providing, by the computing system, the latent sample data to a decoder model, the decoder model configured to decode the latent sample data into a motion forecast including one or more predicted trajectories of the one or more actor features. The computer-implemented method can include receiving, by the computing system, the motion forecast including one or more predicted trajectories of the one or more actor features from the decoder model.

(5) Another example aspect of the present disclosure is directed to a computer-implemented method of training a motion forecasting system. The computer-implemented method can include obtaining, by a computing system including one or more computing devices, a training dataset including one or more training examples labeled with ground truth data, the one or more training examples including one or more actor features and the ground truth data including a ground truth context of the one or more actor features. The computer-implemented method can include providing, by the computing system, the one or more training examples labeled with ground truth data to a latent encoder model, the latent encoder model configured to produce a first latent distribution in response to receipt of the one or more training examples and the ground truth data. The computer-implemented method can include providing, by the computing system, the one or more training examples to a latent prior model, the latent prior model configured to produce a second latent distribution in response to receipt of the one or more training examples. The computer-implemented method can include determining, by the computing system, a training loss based at least in part on the first latent distribution and the second latent distribution. The computer-implemented method can include backpropagating, by the computing system, the training loss through at least the latent prior model to train at least the latent prior model.

(6) Another example aspect of the present disclosure is directed to a computer-implemented method of operating an autonomous vehicle. The computer-implemented method can include obtaining, by the computing system, one or more scene observations. The computer-implemented method can include providing, by the computing system, the one or more scene observations to a feature extraction model, the feature extraction model configured to produce scene data including one or more actor features from the one or more scene observations. The computer-implemented method can include receiving, by the computing system, the scene data including one or more actor features from the feature extraction model. The computer-implemented method can include providing, by the computing system, the scene data to a latent prior model, the latent prior model configured to generate scene latent data in response to receipt of scene data, the scene latent data including one or more latent variables. The computer-implemented method can include obtaining, by the computing system, the scene latent data from the latent prior model. The computer-implemented method can include sampling, by the computing system, one or more latent samples from the scene latent data. The computer-implemented method can include providing, by the computing system, the one or more latent samples to a decoder model, the decoder model configured to decode the latent samples into a motion forecast including one or more predicted trajectories of the one or more actor features. The computer-implemented method can include obtaining, by the computing system, one or more motion forecasts including one or more predicted

trajectories of the one or more actor features from the decoder model. The computer-implemented method can include providing, by the computing system, the one or more predicted trajectories to a motion planning model configured to generate a motion plan for an autonomous vehicle based at least in part on the one or more predicted trajectories. The computer-implemented method can include implementing, by the computing system, the motion plan to control the autonomous vehicle.

(7) Other aspects of the present disclosure are directed to various systems, apparatuses, non-transitory computer-readable media, user interfaces, and electronic devices.

(8) These and other features, aspects, and advantages of various embodiments of the present disclosure will become better understood with reference to the following description and appended claims. The accompanying drawings, which are incorporated in and constitute a part of this specification, illustrate example embodiments of the present disclosure and, together with the description, serve to explain the related principles.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Detailed discussion of embodiments directed to one of ordinary skill in the art is set forth in the specification, which refers to the appended figures, in which:

(2) FIG. 1 depicts a block diagram of an example system for controlling and communicating with a vehicle according to example aspects of the present disclosure.

(3) FIG. 2 depicts an example service infrastructure according to example aspects of the present disclosure.

(4) FIG. 3 depicts a flow diagram of an example feature extraction according to example aspects of the present disclosure.

(5) FIG. 4 depicts a diagram of an example latent distribution according to example aspects of the present disclosure.

(6) FIG. 5 depicts a flow diagram of an example implicit latent variable model according to example aspects of the present disclosure.

(7) FIG. 6 depicts a flow diagram of training an example implicit latent variable model according to example aspects of the present disclosure.

(8) FIG. 7 depicts a flow chart diagram of an example method for determining scene-consistent motion forecasts from sensor data according to example aspects of the present disclosure.

(9) FIG. 8 depicts a flow chart diagram of an example method for recognizing actor features from sensor data according to example aspects of the present disclosure.

(10) FIG. 9 depicts a flow chart diagram of an example method for recognizing actor features from scene features according to example aspects of the present disclosure.

(11) FIG. 10 depicts a flow chart diagram of an example method for training a motion forecasting system according to example aspects of the present disclosure.

(12) FIG. 11 depicts a flow chart diagram of an example method for controlling an autonomous vehicle according to example aspects of the present disclosure.

(13) FIG. 12 depicts a block diagram of a first example computing system according to example aspects of the present disclosure.

(14) FIG. 13 depicts a block diagram of a second example computing system according to example aspects of the present disclosure.

DETAILED DESCRIPTION

(15) Example aspects of the present disclosure are generally directed to systems and methods for latent distribution modeling for scene-consistent motion forecasting. The systems and methods described herein can model interaction between traffic participants, or actors, to provide scene-

consistent motion forecasts of a scene. The motion forecasts can be used in motion planning for controlling an autonomous vehicle. In particular, the systems and methods described herein can characterize joint distributions over future trajectories of one or more traffic participants in a scene by learning a distributed latent representation of the scene. Systems and methods according to example aspects of the present disclosure can model interaction in a scene latent distribution that captures some or all sources of uncertainty. Additionally, systems and methods according to example aspects of the present disclosure can use a deterministic decoder to characterize an implicit joint distribution over actors' future trajectories without any independence assumptions at the output level. This can provide efficient parallel sampling, high expressivity and/or trajectory samples that are substantially more consistent across actors. Samples of the latent representation provided to a deterministic decoder can produce trajectory samples that are consistent across traffic participants and achieve improved interaction understanding. An actor's trajectory can include waypoints over time in the coordinate frame defined by the actor's current position and heading. Thus, systems and methods according to example aspects of the present disclosure can provide for motion plans having improved user comfort and/or safety.

(16) Autonomous vehicles or self-driving vehicles can autonomously transport people and/or goods, providing a safer and/or more efficient solution to transportation. One critical component for autonomous driving is the ability to perceive the world and forecast possible future instantiations of the scene. Producing multi-modal motion forecasts that precisely capture multiple plausible futures consistently for many actors in the scene can present challenges, due at least to the complexity presented by interactions between actors. This complexity can be significant. For instance, the future can be uncertain as actor behaviors may be influenced not only by their own individual goals and intentions but also by the other actors' actions. For instance, an actor at an intersection may choose to turn right or go straight due to its own destination, and yield or go if the behavior of a nearby traffic participant is aggressive or conservative. Moreover, unobserved traffic rules such as the future traffic light states may heavily affect the traffic. Some or all of this information may not be directly observable and thus can require complex reasoning about the scene as a whole, including, for example, its geometry, topology and the interaction between multiple actors.

(17) In a motion planning system, detections and motion forecasts for other actors in the scene may be passed as obstacles to a motion-planner in order to plan a safe maneuver. The distribution over future trajectories may desirably cover the ground-truth for the plan to be safe, but also may desirably exhibit low enough entropy such that a comfortable ride with reasonable progress is achieved. Thus in complex urban environments, an autonomous vehicle can desirably reason about multiple futures separately and plan proactively by understanding how its own actions might influence other actors' behaviors. In addition, in (e.g., closed-loop) self-driving simulators, smart-actor models may be responsible for generating stochastic joint behaviors that are realistic at a scene-level, with actors obeying to underlying scene dynamics with complex interactions.

(18) Thus, it can be desirable to learn a joint distribution over actors' future trajectories that characterizes how the scene might unroll as a whole. A joint distribution over actors' future trajectories can provide for samples that are socially consistent across an entire scene, which can provide a motion planner system with improved reasoning about possible future scenarios separately. This can also provide for realistic simulation of complex traffic scenes. However, this may be intractable for some existing systems and methods, especially for complex scenes. To solve this problem, some existing motion forecasting approaches assume marginal independence across actors' future trajectories. This can cause systems to fail to achieve scene-consistent futures. Alternatively, auto-regressive formulations model interactions at the output level, but require sequential sampling which can result in slow inference and compounding errors.

(19) Various factorizations of the joint distribution over actors' trajectories of independence assumptions have been proposed to sidestep the intractability of true conditional dependence. One

simple approximation is to assume independent futures across actors and time steps. Some approaches directly regress the parameters of a mixture of Gaussians over time, which provides efficient sampling but can suffer from low expressivity and unstable optimization. Non-parametric approaches have also been proposed to characterize the multi-modality of one actor's individual behavior. For instance, some approaches score trajectory samples from a finite set with limited coverage. Some other approaches predict an occupancy grid at different future horizons, which can be very memory consuming. Some other approaches propose to learn a one-step policy that predicts the next waypoint based on the previous history, avoiding the time independence assumption. Variational methods have also been proposed to learn an actor independent latent distribution to capture unobserved actor dynamics such as goals. However, none of these existing methods can accurately characterize the joint distribution in interactive situations, since the generative process is independent per actor.

(20) Another existing approach to characterize the behavior of multiple actors jointly is autoregressive generation with social mechanisms, which predict the distribution over the next trajectory waypoint of each actor conditioned on the previous states of all actors. Autoregressive approaches, however, can suffer from compounding errors. For instance, during training, the model is fed the ground-truth while during inference, the model must rely on approximate samples from the learned distribution. The objective function underlying this method pushes the conditional distributions to model the marginal distributions instead. Moreover, these methods require sequential sampling, which may not be amenable to some real-time applications such as self-driving. Furthermore, capturing uncertainty and multi-modality at the actor level may not guarantee that if samples are taken from each of the actors independently, the samples will be consistent with each other.

(21) For instance, consider an example where two actors approach an intersection. Assuming they have similar speed, then an independent output for each actor may be similar. These marginals could be accurate characterizations of the world when using a simple model, since the two prominent modes at the scene-level is that one actor yields and the other one goes, or the other way around. However, this model may fail to provide scene consistent samples. Since the output distributions for each actor are independent, when a sample is obtained from each of them, the samples may describe inconsistent futures, such as a future where both actors go, resulting in a collision.

(22) Systems and methods according to example aspects of the present disclosure can provide solutions to these and other challenges. For instance, systems and methods according to example aspects of the present disclosure can characterize a joint distribution over motion forecasts via an implicit latent variable model (ILVM). The implicit latent variable model can model a latent distribution that can summarize unobserved scene dynamics given input sensor data and/or scene features. This can be challenging given that modern roads may present very complex geometries and topologies that can make every intersection unique. Furthermore, this can be challenging given that the dynamic environment of the scene may be only partially observed through sensor returns. Finally, challenges can be encountered as the number of actors in a scene is variable.

(23) To address these and/or other challenges, systems and methods according to example aspects of the present disclosure can model a scene as an interaction graph including one or more nodes. For instance, the nodes can correspond to traffic participants or actors (e.g., actor features). This interaction graph can be used to produce a scene latent distribution of one or more latent variables. The scene latent distribution can be partitioned into a distributed representation among actors. For instance, scene interaction modules including, for example, graph neural networks (GNN) can be used to encode the full scene into the scene latent distribution and/or to decode latent samples from the scene latent distribution into socially consistent future trajectories. For instance, a deterministic decoder can frame the decoding of all actors' trajectories as a deterministic mapping from the inputs and scene latent samples. This can provide that the latent variables capture all the

stochasticity in the generative process. This can also provide for efficient multi-sample inference via parallel sampling.

(24) For instance, example aspects of the present disclosure are directed to a computer-implemented method for determining scene-consistent motion forecasts from sensor data. The computer-implemented method can be implemented by any suitable computing system, such as an autonomous vehicle navigation system.

(25) The computer-implemented method can include obtaining (e.g., by a computing system including one or more computing devices) scene data including one or more actor features. In some implementations, the scene data can be extracted or otherwise obtained from one or more scene observations. For instance, the method can include obtaining (e.g., by the computing system) one or more scene observations. The scene observations can be or can include data descriptive of sensor observations from one or more sensors configured to observe the scene, such as, for example, one or more sensors mounted on an autonomous vehicle. The sensors can be any suitable sensors, such as, for example, cameras, LIDAR sensors, etc. As one example, in some implementations, the scene observations can be and/or can include a three-dimensional (3D) LIDAR point cloud. In some implementations, the LIDAR data can be voxelized. In some implementations, the scene observations can be oriented in a “Birds-Eye View” (BEV) or top-down representation.

(26) In some implementations, the scene observations can additionally include map data, such as data descriptive of properties of roads, crosswalks, signage, intersections, railroads, buildings, and/or other terrain features of the scene. In some implementations, the map data can be rasterized. The map data can encode traffic elements such as intersections, lanes, roads, and traffic lights. In some implementations, elements with different semantics are encoded into different channels in the raster. Map elements that are rasterized can be or can include, for example, drivable surface polygons, road polygons, intersection polygons, straight vehicle lane polygons, dedicated left and right vehicle lane polygons, dedicated bike lane polygons, dedicated bus lane polygons, centerline markers for all lanes, lane dividers for all lanes with semantics (e.g., allowed to cross, not allowed to cross, might be allowed to cross).

(27) In some implementations, the height dimension of the sensor observations can be normalized with dense ground-height information provided by map data. In some implementations, multiple LiDAR sweeps can be used to exploit motion cues by compensating the ego-motion, such as by projecting the past sweeps to the coordinate frame of the current sweep). Furthermore, in some implementations, the height and time dimension is raveled into the channel dimension, to provide for the use of 2D convolution to process spatial and/or temporal information efficiently. The final representation may thus be a 3D occupancy tensor.

(28) The method can further include providing (e.g., by the computing system) the one or more scene observations to a feature extraction model. The feature extraction model can include one or more neural networks configured to produce scene data including one or more actor features from the one or more scene observations. For instance, in some cases, the features can be extracted from raw sensor data and HD maps in a differentiable manner, such that perception and motion forecasting can be trained jointly end-to-end. In some implementations, the feature extraction model can be or can include a backbone network. For instance, the backbone network can be a lightweight backbone network adapted for feature extraction. In some implementations, two separate streams can be instantiated such that the voxelized LiDAR and rasterized map are processed separately. The resulting features from both streams can then be concatenated feature-wise (e.g., if they share the same spatial resolution) and fused by a convolutional header. These extracted features can inform both the downstream detection and motion forecasting networks. The method can then include receiving (e.g., by the computing system) the scene data including one or more actor features from the feature extraction model.

(29) In some implementations, the feature extraction model can include a scene feature extraction model and/or an actor feature recognition model. For instance, the method can include providing

(e.g., by the computing system) the one or more scene observations to a scene feature extraction model. The scene feature extraction model can include one or more neural networks configured to extract one or more scene features from the one or more scene observations. The method can then include receiving (e.g., by the computing system) the one or more scene features from the scene feature extraction model. The scene features may be features that may not each correspond to an actor in the scene, such as a global feature map. For instance, a CNN-based perception backbone network architecture can be used to extract rich geometrical and motion features about the whole scene from a past history of voxelized LiDAR point clouds and/or a raster map.

(30) The method can then include providing (e.g., by the computing system) the one or more scene features to an actor feature recognition model. The actor feature recognition model can parse the scene features into actor features that correspond to an actor. For instance, the one or more actor features can include data descriptive of an actor context of one or more traffic participants. For instance, the actor feature recognition model can be configured to extract spatial feature maps for bounding boxes from the one or more scene features by rotated region of interest (ROI) align. Rotated ROI align can be applied to extract (e.g., fixed size) spatial feature maps for bounding boxes with arbitrary shapes and rotations from the scene features (e.g., the global feature map extracted by the backbone). For instance, rotated ROI align can provide actor contexts for each actor.

(31) Additionally and/or alternatively, the actor feature recognition model can be configured to pool a region around each spatial feature map to produce pooled actor features. For instance, a region around each actor in its frame can be pooled, such as with an axis defined by the actor's centroid orientation. The pooled actor features may be larger than the eventual actor features.

(32) Additionally and/or alternatively, the actor feature recognition model can be configured to downsample the pooled actor features by applying one or more downsampling convolutional neural networks. As one example, a 4-layer down-sampling convolutional network can be applied.

(33) Additionally and/or alternatively, the actor feature recognition model can be configured to max-pool along spatial dimensions to reduce each pooled actor feature to a respective actor feature of the one or more actor features. For instance, after the downsampling CNN is applied, it can be followed by max-pooling along the spatial dimensions to reduce the feature map to a desired dimensional feature vector per actor. One example convolutional network uses a dilation factor of 2 for the convolutional layers to enlarge the receptive field for the per-actor features, which can improve performance. The method can then include receiving (e.g., by the computing system) the one or more actor features from the actor feature recognition model.

(34) For instance, in some implementations, the (e.g., scene) feature extraction model (e.g., the backbone model) includes two convolutional layers to output a classification or confidence score and/or a bounding box for each anchor location (e.g., each scene feature). These features are eventually reduced to the final set of candidates by applying non-maximal suppression (NMS) and finally thresholding low probability detections (given by the desired common recall). In some implementations, a backbone network along with features for object detection and per actor feature extraction are provided. The proposed mixture of trajectories output parameterization, where each way-point is a gaussian, is then used. In some cases, these baselines may not obtain temporally consistent samples, since the gaussians are independent across time (e.g., the models are not auto-regressive). To solve this, a heuristic sampler can be used to obtain temporally consistent samples from this model. The sampled trajectories are extracted using the re-parameterization technique for a bi-variate normal, where the model predicts a normal distribution per waypoint.

(35) In some cases, the noise can be constant across time for a given sample and actor. Intuitively, having a constant noise across time steps can provide sampled waypoints whose relative location with respect to its predicted mean and covariance is constant across time (e.g., translated by the predicted mean and scaled by the predicted covariance per time). In some cases, to address the compounding error problem found in some auto-regressive models, adjustments can be made to the

training procedure to account for the noise in the conditioning space. To help simulate the noise it sees during inference, gaussian noise can be added to the conditioning state. The amount of noise expected between time-steps can be tuned.

(36) In some implementations, an object detection module can be responsible for recognizing other traffic participants in a scene, followed by a motion forecasting module that predicts how the scene might unroll given the current state or actor state of each actor. The actor state may be a compact representation of an actor, including qualities such as pose, velocity, and acceleration. This can be beneficial in some cases; however, it can be difficult to incorporate uncertainty due to sensor noise or occlusion. In some implementations, these two tasks can be combined by having a single model (e.g., a single fully convolutional backbone network) predict both the current and future states of actors. For instance, a single fully convolutional backbone network can predict both the current state and future state(s) for each pixel (e.g., in a bird's eye view grid) directly from sensor data, such as a voxelized LiDAR point-cloud, and/or map data, such as a semantic raster of an HD map. This approach can propagate uncertainty between the two tasks in the feature space, without the need of explicit intermediate representations.

(37) For instance, the perception and prediction tasks seek to understand where other actors are currently located and/or how they might move in the next few seconds. This can be accomplished by leveraging (e.g., 3D) sensor data such as LiDAR point clouds for dynamic object recognition and/or high-definition maps which provide information about the static part of the environment. For instance, scene features can be extracted from LiDAR and maps and subsequently fused to produce rich features. Once rich features from the whole scene are extracted, object detection can be performed to recognize actor features corresponding to actors in the scenes.

(38) Additionally and/or alternatively, the method can include providing (e.g., by the computing system) the scene data to a latent prior model. The latent prior model can be configured to generate scene latent data in response to receipt of scene data. In some implementations, the latent prior model can be or can include a scene interaction module including one or more graph neural networks. The scene latent data can be or can include one or more latent variables. In some implementations, the scene latent data can include a latent distribution that is partitioned into one or more latent variables. For instance, in some implementations, the one or more latent variables can be respective to the one or more actor features such that each actor feature has an associated latent variable of the scene latent data that is anchored to the actor feature. In some implementations, the one or more latent variables can be or can include one or more continuous latent variables. Additionally and/or alternatively, the method can include obtaining (e.g., by the computing system) the scene latent data from the latent prior model.

(39) For instance, the generative process of future trajectories over actors can be formulated with a latent variable model including one or more latent variables that intuitively capture unobserved scene dynamics such as actor goals and style, multi-agent interactions, or future traffic light states. This modeling intuitively encourages the latent distribution to capture stochasticity in the generative process. In some implementations, the latent distribution can be a continuous latent distribution including one or more continuous latent variables for high expressivity.

(40) Producing a latent distribution that can capture all the uncertainties in any scenario can present challenges, as scenarios can vary drastically in the number of actors, the road topology, and/or traffic rules. This challenge can be mitigated by partitioning the scene latent distribution and obtaining a distributed representation where each latent variable in the scene latent distribution is anchored to a respective actor, such as anchored to a particular node in an interaction graph with traffic participants as nodes. The distributed representation may provide the benefit of naturally scaling the capacity of the latent distribution as the number of actors grows. Furthermore, the anchoring may give the model an inductive bias that eases the learning of a scene latent distribution. Intuitively, each anchored latent variable encodes unobserved dynamics most relevant to its respective actor, including interactions with neighboring actors and traffic rules that apply in

its locality. For instance, each latent variable can be represented as a diagonal multivariate Gaussian. Each latent variable can be conditioned on all actors such that the latent distribution is not marginally independent across actors, although factorized. For instance, despite anchoring each partition of the scene latent to an actor, each individual latent variable can contain information about the full scene, since each final node representation is dependent on the entire input because of the message propagation in the fully-connected interaction graph.

(41) Additionally and/or alternatively, the method can include sampling (e.g., by the computing system) latent sample data from the scene latent data. For instance, the scene latent data (e.g., the one or more latent variables) can define a latent (e.g., a latent distribution) that can be sampled to produce latent samples of the scene latent data. The latent sample data can define a possible future for the actors (e.g., the actor features).

(42) Additionally and/or alternatively, the method can include providing (e.g., by the computing system) the latent sample data to a decoder model. The decoder model can be configured to decode the latent sample data into a motion forecast including one or more predicted trajectories of the one or more actor features. For instance, the decoder model can produce a motion forecast from latent samples. Additionally and/or alternatively, the method can include receiving (e.g., by the computing system) the motion forecast including one or more predicted trajectories of the one or more actor features from the decoder model.

(43) In some implementations, the decoder model can be or can include a deterministic decoder model. For instance, the decoder can be or can include a deterministic mapping to implicitly characterize the joint probability, such as opposed to explicitly representing it in a parametric form. This approach can provide for evaluation without factorizing the joint distribution and thus can sidestep the associated challenges. In particular, the deterministic decoder can be highly beneficial for generating socially consistent trajectories. In this framework, generating scene-consistent future trajectories across actors is simple and highly efficient, as it may be performed with only one stage of parallel sampling.

(44) In some implementations, the decoder model can be or can include a scene interaction module including one or more graph neural networks. For instance, the decoder including a scene interaction module can predict a realization of the future at the scene level via message parsing. As an example, each actor trajectory can be established with respect to samples from each latent variable of the scene latent data and/or each actor feature. This can provide for improved reasoning about multi-agent interactions such as car following, yielding, etc. For instance, each actor context can be initialized as a node in the decoder SIM. After a round of message parsing, each node can then contain an updated representation of a respective actor that takes into account the underlying dynamics of the scene summarized in the latent distribution. Finally, the trajectory sample for the actor can be deterministically decoded, such as by the output function of the SIM. For instance, the output function can be the deterministic mapping. This can be performed without requiring any additional sampling steps. The trajectory-level scene sample can thus be a collection of all actor trajectories.

(45) In some implementations, the decoder model can be or can include a specified and tractable conditional likelihood. In this implementation, many tools are available for inference and learning. As one example, variational inference, such as the variational auto-encoder (VAE), can be used.

(46) In some implementations, the decoder can be characterized via a stochastic sampling procedure where a conditional dependence is not specified. In this case, likelihood-free inference methods can be required for learning. As examples, density estimation by comparison using either density ratio (GAN) or density difference (MMD). These methods, however, may be more difficult to optimize.

(47) In some implementations the prior model, the encoder model, and/or the decoder model can include a scene interaction model. The scene interaction model can be configured to model the latent distribution as an interaction graph including one or more nodes representative of the one or

more actor features and one or more edges representative of interactions between the one or more actor features. For instance, in some implementations, the scene interaction model can include one or more graph neural networks. In some implementations, a message function of the one or more graph neural networks can include a multi-layer perceptron model that takes as input one or more terminal nodes of the one or more nodes at a previous propagation step of the one or more graph neural networks. For instance, the edge or message function of the graph neural network(s) in the scene interaction module can include, for example, a 3-layer MLP (multi-layer perceptron) that takes as input the hidden states of the two terminal nodes at each edge in the graph at the previous propagation step. Additionally and/or alternatively, the input can include the projected coordinates of their corresponding bounding boxes. In some implementations, an aggregation function of the one or more graph neural networks includes a feature-wise max-pooling aggregation function. In some implementations, a gated recurrent unit cell is configured to update a state (e.g., hidden state) of the one or more nodes. In some implementations, the scene interaction module can include an output network that outputs the results from the graph propagations, such as a 2-layer MLP.

(48) The scene interaction module can model the latent distribution as an interaction graph, which can provide improved understanding of spatial information. This spatial information can be beneficial in jointly forecasting future trajectories of each actor. For instance, the node state of each node can be initialized with a set of actor features and known spatial information. The spatial information can include information such as relative coordinates of the actors relative to their peers or neighbors. In some cases, during object detection and local feature extraction around each actor, however, it may not be possible to include some necessary global information due to the limited receptive field and/or the translation invariance of convolutional neural networks. To remedy this, the node states can be initialized as the concatenation of the deep local features and the spatial information of each actor or node in the graph, such as its location, heading and/or its dimensions (e.g., in Bird's Eye View). A learned double edge function can propagate messages around the nodes in the graph. Given these messages, each actor can aggregate the messages (e.g., via max pooling) to update a respective node state. In some implementations, the scene interaction model can perform a single round of message passing to update the nodes' representation, taking into account spatiotemporal relationships. The scene interaction module in the prior, encoder and/or decoder can capture scene-level understanding that is not present with independence assumptions at the latent or output level.

(49) The systems and methods described herein can provide for efficiently sampling multiple possible futures from the latent distribution of the scene latent distribution. For instance, in some implementations, the method can include sampling (e.g., by the computing system) second latent sample data from the scene latent data. The second latent sample data can be descriptive of a differing possible future from the latent sample data. Additionally and/or alternatively, the method can include providing (e.g., by the computing system) the second latent sample data to the decoder model and receiving (e.g., by the computing system) a second motion forecast including one or more second predicted trajectories of the one or more actor features from the decoder model. The second predicted trajectories can differ from the predicted trajectories of other samples.

(50) For instance, a first sample from the scene latent distribution provided to the decoder can produce one possible realization of the future trajectories. A second sample can result in a distinct future. The sampling process is illustrated sequentially for the purposes of illustration, parallel sampling and/or decoding can be employed in accordance with example aspects of the present disclosure. For instance, the samples may be independent, as the stochasticity in the system is present in the latent distribution (e.g., as opposed to the decoder).

(51) In other implementations, example aspects of the present disclosure are directed to a computer-implemented method of training a motion forecasting system. For instance, the method can include training an implicit latent variable model according to example aspects of the present disclosure. In some implementations, the model can be fully differentiable and can be trained end-

to-end through backpropagation using a multi-task objective.

(52) The method can include obtaining (e.g., by a computing system including one or more computing devices) a training dataset including one or more training examples labeled with ground truth data. For instance, the one or more training examples can include one or more actor features. Additionally and/or alternatively, the ground truth data can include a ground truth context of the one or more actor features. The ground truth context can be descriptive of a known context of the actor features, such as a known position, pose, velocity, etc.

(53) The method can include providing (e.g., by the computing system) the one or more training examples labeled with ground truth data to a latent encoder model. The latent encoder model can be configured to produce a first latent distribution in response to receipt of the one or more training examples and the ground truth data. For instance, in some cases, integration over the scene latent distribution is intractable, in which case amortized variational inference can be used. For instance, by introducing an encoder distribution to approximate the true posterior, the learning problem can be reformulated as a maximization of the Evidence Lower BOund (ELBO). In some implementations, the latent encoder model can include a scene interaction module. For instance, after running one round of message passing, the scene interaction module can predict the distribution over latent variables.

(54) The encoder model can approximate the true posterior latent distribution. This model may also be called a recognition network. Because the encoder model receives the ground truth data (e.g., the target output) as an input, it can recognize scene dynamics that are unobserved by the latent prior model. In this manner, the encoder may only be used during training, since it requires access to the ground-truth future trajectories. For instance, the encoder may be omitted from deployed models and/or included at an online training system.

(55) The method can include providing (e.g., by the computing system) the one or more training examples to a latent prior model. The latent prior model can be configured to produce a second latent distribution in response to receipt of the one or more training examples. For instance, the latent prior model can be agnostic to the ground truth data such that the latent prior model is usable during inference (e.g., when ground truth data is unavailable).

(56) The method can include determining (e.g., by the computing system) a training loss based at least in part on the first latent distribution and the second latent distribution. The training loss can be or can include various losses from components of the model. For instance, in some implementations, the training loss can be or can include a KL divergence loss between the first latent distribution and the second latent distribution.

(57) Additionally and/or alternatively, in some implementations, the training loss can include a feature loss. The feature loss can be a loss in recognized features and/or actors. For instance, the feature loss can be determined by providing (e.g., by the computing system) training scene observations to a feature extraction model. The feature extraction model can be or can include a scene feature extraction model and/or an actor feature recognition model. Determining the feature loss can further include receiving (e.g., by the computing system) one or more predicted features from the feature extraction model. Determining the feature loss can further include determining (e.g., by the computing system) a feature loss between the one or more predicted features and the ground truth data. The feature loss can include a cross-entropy loss between the one or more predicted features and one or more training features of the ground truth data. Additionally and/or alternatively, the feature loss can include a regression loss between bounding boxes of the one or more predicted features and bounding boxes of the ground truth data. For instance, in some implementations, the feature loss includes a cross entropy loss applied to detection scores of the predicted features and/or a smooth L1 applied to the bounding box regression targets of the one or more predicted features. The method can include backpropagating (e.g., by the computing system) the training loss through at least the latent prior model to train at least the latent prior model. For instance, the training loss can be backpropagated through some or all trainable portions of the

implicit latent variable model, such as an object detection or feature extraction model, latent prior model, decoder model, encoder model, etc.

(58) In some implementations, the training loss can include a reconstruction loss. For instance, the reconstruction loss can be obtained by sampling (e.g., by the computing system) the first latent distribution to produce one or more first latent samples. The first latent distribution can include first sample data, such as a first set of samples from the latent distribution. Determining the reconstruction loss can further include sampling (e.g., by the computing system) the second latent distribution to produce one or more second latent samples. Determining the reconstruction loss can further include providing (e.g., by the computing system) the one or more first latent samples to a decoder model and receiving (e.g., by the computing system) one or more first predicted trajectories from the decoder model. The first predicted trajectories can be representative of a first possible future of the actor features. Determining the reconstruction loss can further include providing (e.g., by the computing system) the one or more second latent samples to the decoder model and receiving (e.g., by the computing system) one or more second predicted trajectories from the decoder model. The second predicted trajectories can be representative of a second possible future of the actor features that desirably matches the first possible future. Determining the reconstruction loss can include determining (e.g., by the computing system) a reconstruction loss based at least in part on the one or more first predicted trajectories and the one or more second predicted trajectories. For instance, the reconstruction loss can be based at least in part on a difference between the one or more first predicted trajectories and the one or more second predicted trajectories. In some implementations, the reconstruction loss can include a Huber loss.

(59) The motion forecasts produced as described according to example aspects of the present disclosure can be useful for motion planning and/or operating an autonomous vehicle. For instance, example aspects of the present disclosure are directed to a computer-implemented method of operating an autonomous vehicle. The systems and methods described herein can be used to generate a motion plan for safely navigating an autonomous vehicle.

(60) The method can include obtaining one or more predicted trajectories of one or more actors. For instance, the one or more predicted trajectories can be obtained from an implicit latent variable model according to example aspects of the present disclosure.

(61) Additionally and/or alternatively, the method can include providing (e.g., by the computing system) the one or more predicted trajectories to a motion planning model configured to generate a motion plan for an autonomous vehicle based at least in part on the one or more predicted trajectories. For instance, the motion planning model can generate a motion plan that safely navigates the scene in view of the predicted trajectories of the actors in the scene.

(62) Additionally and/or alternatively, the method can include implementing (e.g., by the computing system) the motion plan to control the autonomous vehicle. For instance, one or more control devices (e.g., throttle, brakes, steering, etc.) can be controlled based at least in part on the motion plan to implement the motion plan and thus navigate the autonomous vehicle.

(63) Systems and methods according to example aspects of the present disclosure can achieve a number of technical effects and benefits, including but not limited to improvements to computing technologies. For instance, systems and methods (e.g., including an implicit latent variable model) can achieve improved generation of scene-consistent samples while producing less entropic joint distributions that also better cover the ground-truth. Moreover, when using the scene-consistent motion forecasts, a motion planner (e.g., including a conventional motion planner) can plan safer and more comfortable trajectories.

(64) Referring now to the FIGS., example aspects of the present disclosure will be discussed in more detail. FIG. 1 depicts a block diagram of an example system **100** for controlling and communicating with a vehicle according to example aspects of the present disclosure. As illustrated, FIG. 1 shows a system **100** that can include a vehicle **105** and a vehicle computing system **110** associated with the vehicle **105**. The vehicle computing system **110** can be located

onboard the vehicle **105** (e.g., it can be included on and/or within the vehicle **105**).

(65) The vehicle **105** incorporating the vehicle computing system **110** can be various types of vehicles. For instance, the vehicle **105** can be an autonomous vehicle. The vehicle **105** can be a ground-based autonomous vehicle (e.g., car, truck, bus, etc.). The vehicle **105** can be an air-based autonomous vehicle (e.g., airplane, helicopter, vertical take-off and lift (VTOL) aircraft, etc.). The vehicle **105** can be a light weight elective vehicle (e.g., bicycle, scooter, etc.). The vehicle **105** can be another type of vehicle (e.g., watercraft, etc.). The vehicle **105** can drive, navigate, operate, etc. with minimal and/or no interaction from a human operator (e.g., driver, pilot, etc.). In some implementations, a human operator can be omitted from the vehicle **105** (and/or also omitted from remote control of the vehicle **105**). In some implementations, a human operator can be included in the vehicle **105**.

(66) The vehicle **105** can be configured to operate in a plurality of operating modes. The vehicle **105** can be configured to operate in a fully autonomous (e.g., self-driving) operating mode in which the vehicle **105** is controllable without user input (e.g., can drive and navigate with no input from a human operator present in the vehicle **105** and/or remote from the vehicle **105**). The vehicle **105** can operate in a semi-autonomous operating mode in which the vehicle **105** can operate with some input from a human operator present in the vehicle **105** (and/or a human operator that is remote from the vehicle **105**). The vehicle **105** can enter into a manual operating mode in which the vehicle **105** is fully controllable by a human operator (e.g., human driver, pilot, etc.) and can be prohibited and/or disabled (e.g., temporary, permanently, etc.) from performing autonomous navigation (e.g., autonomous driving, flying, etc.). The vehicle **105** can be configured to operate in other modes such as, for example, park and/or sleep modes (e.g., for use between tasks/actions such as waiting to provide a vehicle service, recharging, etc.). In some implementations, the vehicle **105** can implement vehicle operating assistance technology (e.g., collision mitigation system, power assist steering, etc.), for example, to help assist the human operator of the vehicle **105** (e.g., while in a manual mode, etc.).

(67) To help maintain and switch between operating modes, the vehicle computing system **110** can store data indicative of the operating modes of the vehicle **105** in a memory onboard the vehicle **105**. For example, the operating modes can be defined by an operating mode data structure (e.g., rule, list, table, etc.) that indicates one or more operating parameters for the vehicle **105**, while in the particular operating mode. For example, an operating mode data structure can indicate that the vehicle **105** is to autonomously plan its motion when in the fully autonomous operating mode. The vehicle computing system **110** can access the memory when implementing an operating mode.

(68) The operating mode of the vehicle **105** can be adjusted in a variety of manners. For example, the operating mode of the vehicle **105** can be selected remotely, off-board the vehicle **105**. For example, a remote computing system (e.g., of a vehicle provider and/or service entity associated with the vehicle **105**) can communicate data to the vehicle **105** instructing the vehicle **105** to enter into, exit from, maintain, etc. an operating mode. By way of example, such data can instruct the vehicle **105** to enter into the fully autonomous operating mode.

(69) In some implementations, the operating mode of the vehicle **105** can be set onboard and/or near the vehicle **105**. For example, the vehicle computing system **110** can automatically determine when and where the vehicle **105** is to enter, change, maintain, etc. a particular operating mode (e.g., without user input). Additionally, or alternatively, the operating mode of the vehicle **105** can be manually selected via one or more interfaces located onboard the vehicle **105** (e.g., key switch, button, etc.) and/or associated with a computing device proximate to the vehicle **105** (e.g., a tablet operated by authorized personnel located near the vehicle **105**). In some implementations, the operating mode of the vehicle **105** can be adjusted by manipulating a series of interfaces in a particular order to cause the vehicle **105** to enter into a particular operating mode.

(70) The vehicle computing system **110** can include one or more computing devices located onboard the vehicle **105**. For example, the computing device(s) can be located on and/or within the

vehicle **105**. The computing device(s) can include various components for performing various operations and functions. For instance, the computing device(s) can include one or more processors and one or more tangible, non-transitory, computer readable media (e.g., memory devices, etc.). The one or more tangible, non-transitory, computer readable media can store instructions that when executed by the one or more processors cause the vehicle **105** (e.g., its computing system, one or more processors, etc.) to perform operations and functions, such as those described herein for controlling an autonomous vehicle, communicating with other computing systems, etc.

(71) The vehicle **105** can include a communications system **115** configured to allow the vehicle computing system **110** (and its computing device(s)) to communicate with other computing devices. The communications system **115** can include any suitable components for interfacing with one or more network(s) **120**, including, for example, transmitters, receivers, ports, controllers, antennas, and/or other suitable components that can help facilitate communication. In some implementations, the communications system **115** can include a plurality of components (e.g., antennas, transmitters, and/or receivers) that allow it to implement and utilize multiple-input, multiple-output (MIMO) technology and communication techniques.

(72) The vehicle computing system **110** can use the communications system **115** to communicate with one or more computing device(s) that are remote from the vehicle **105** over one or more networks **120** (e.g., via one or more wireless signal connections). The network(s) **120** can exchange (send and/or receive) signals (e.g., electronic signals), data (e.g., data from a computing device), and/or other information and include any combination of various wired (e.g., twisted pair cable) and/or wireless communication mechanisms (e.g., cellular, wireless, satellite, microwave, and radio frequency) and/or any desired network topology (or topologies). For example, the network(s) **120** can include a local area network (e.g., intranet), wide area network (e.g., Internet), wireless LAN network (e.g., via Wi-Fi), cellular network, a SATCOM network, VHF network, a HF network, a WiMAX based network, and/or any other suitable communication network (or combination thereof) for transmitting data to and/or from the vehicle **105** and/or among computing systems.

(73) In some implementations, the communications system **115** can also be configured to enable the vehicle **105** to communicate with and/or provide and/or receive data and/or signals from a remote computing device associated with a user **125** and/or an item (e.g., an item to be picked-up for a courier service). For example, the communications system **115** can allow the vehicle **105** to locate and/or exchange communications with a user device **130** of a user **125**. In some implementations, the communications system **115** can allow communication among one or more of the system(s) on-board the vehicle **105**.

(74) As shown in FIG. 1, the vehicle **105** can include one or more sensors **135**, an autonomy computing system **140**, a vehicle interface **145**, one or more vehicle control systems **150**, and other systems, as described herein. One or more of these systems can be configured to communicate with one another via one or more communication channels. The communication channel(s) can include one or more data buses (e.g., controller area network (CAN)), on-board diagnostics connector (e.g., OBD-II), and/or a combination of wired and/or wireless communication links. The onboard systems can send and/or receive data, messages, signals, etc. amongst one another via the communication channel(s).

(75) The sensor(s) **135** can be configured to acquire sensor data **155** (e.g., one or more sensor observations). The sensor(s) **135** can be external sensors configured to acquire external sensor data. This can include sensor data associated with the surrounding environment of the vehicle **105**. The surrounding environment of the vehicle **105** can include/be represented in the field of view of the sensor(s) **135**. For instance, the sensor(s) **135** can acquire image and/or other data of the environment outside of the vehicle **105** and within a range and/or field of view of one or more of the sensor(s) **135**. The sensor(s) **135** can include one or more Light Detection and Ranging (LIDAR) systems, one or more Radio Detection and Ranging (RADAR) systems, one or more cameras (e.g., visible spectrum cameras, infrared cameras, etc.), one or more motion sensors, one

or more audio sensors (e.g., microphones, etc.), and/or other types of imaging capture devices and/or sensors. The one or more sensors can be located on various parts of the vehicle **105** including a front side, rear side, left side, right side, top, and/or bottom of the vehicle **105**. The sensor data **155** can include image data (e.g., 2D camera data, video data, etc.), RADAR data, LIDAR data (e.g., 3D point cloud data, etc.), audio data, and/or other types of data. The vehicle **105** can also include other sensors configured to acquire data associated with the vehicle **105**. For example, the vehicle **105** can include inertial measurement unit(s), wheel odometry devices, and/or other sensors.

(76) In some implementations, the sensor(s) **135** can include one or more internal sensors. The internal sensor(s) can be configured to acquire sensor data **155** associated with the interior of the vehicle **105**. For example, the internal sensor(s) can include one or more cameras, one or more infrared sensors, one or more motion sensors, one or more weight sensors (e.g., in a seat, in a trunk, etc.), and/or other types of sensors. The sensor data **155** acquired via the internal sensor(s) can include, for example, image data indicative of a position of a passenger or item located within the interior (e.g., cabin, trunk, etc.) of the vehicle **105**. This information can be used, for example, to ensure the safety of the passenger, to prevent an item from being left by a passenger, confirm the cleanliness of the vehicle **105**, remotely assist a passenger, etc.

(77) In some implementations, the sensor data **155** can be indicative of one or more objects within the surrounding environment of the vehicle **105**. The object(s) can include, for example, vehicles, pedestrians, bicycles, and/or other objects. The object(s) can be located in front of, to the rear of, to the side of, above, below the vehicle **105**, etc. The sensor data **155** can be indicative of locations associated with the object(s) within the surrounding environment of the vehicle **105** at one or more times. The object(s) can be static objects (e.g., not in motion) and/or dynamic objects/actors (e.g., in motion or likely to be in motion) in the vehicle's environment. The sensor(s) **135** can provide the sensor data **155** to the autonomy computing system **140**.

(78) In addition to the sensor data **155**, the autonomy computing system **140** can obtain map data **160**. The map data **160** can provide detailed information about the surrounding environment of the vehicle **105** and/or the geographic area in which the vehicle was, is, and/or will be located. For example, the map data **160** can provide information regarding: the identity and location of different roadways, road segments, buildings, or other items or objects (e.g., lampposts, crosswalks and/or curb); the location and directions of traffic lanes (e.g., the location and direction of a parking lane, a turning lane, a bicycle lane, or other lanes within a particular roadway or other travel way and/or one or more boundary markings associated therewith); traffic control data (e.g., the location and instructions of signage, traffic lights, and/or other traffic control devices); obstruction information (e.g., temporary or permanent blockages, etc.); event data (e.g., road closures/traffic rule alterations due to parades, concerts, sporting events, etc.); nominal vehicle path data (e.g., indicate of an ideal vehicle path such as along the center of a certain lane, etc.); and/or any other map data that provides information that assists the vehicle computing system **110** in processing, analyzing, and perceiving its surrounding environment and its relationship thereto. In some implementations, the map data **160** can include high definition map data. In some implementations, the map data **160** can include sparse map data indicative of a limited number of environmental features (e.g., lane boundaries, etc.). In some implementations, the map data can be limited to geographic area(s) and/or operating domains in which the vehicle **105** (or autonomous vehicles generally) may travel (e.g., due to legal/regulatory constraints, autonomy capabilities, and/or other factors).

(79) The vehicle **105** can include a positioning system **165**. The positioning system **165** can determine a current position of the vehicle **105**. This can help the vehicle **105** localize itself within its environment. The positioning system **165** can be any device or circuitry for analyzing the position of the vehicle **105**. For example, the positioning system **165** can determine position by using one or more of inertial sensors (e.g., inertial measurement unit(s), etc.), a satellite positioning system, based on IP address, by using triangulation and/or proximity to network access points or

other network components (e.g., cellular towers, WiFi access points, etc.) and/or other suitable techniques. The position of the vehicle **105** can be used by various systems of the vehicle computing system **110** and/or provided to a remote computing system. For example, the map data **160** can provide the vehicle **105** relative positions of the elements of a surrounding environment of the vehicle **105**. The vehicle **105** can identify its position within the surrounding environment (e.g., across six axes, etc.) based at least in part on the map data **160**. For example, the vehicle computing system **110** can process the sensor data **155** (e.g., LIDAR data, camera data, etc.) to match it to a map of the surrounding environment to get an understanding of the vehicle's position within that environment. Data indicative of the vehicle's position can be stored, communicated to, and/or otherwise obtained by the autonomy computing system **140**.

(80) The autonomy computing system **140** can perform various functions for autonomously operating the vehicle **105**. For example, the autonomy computing system **140** can perform the following functions: perception **170A**, prediction **170B**, and motion planning **170C**. For example, the autonomy computing system **140** can obtain the sensor data **155** via the sensor(s) **135**, process the sensor data **155** (and/or other data) to perceive its surrounding environment, predict the motion of objects within the surrounding environment, and generate an appropriate motion plan through such surrounding environment. In some implementations, these autonomy functions can be performed by one or more sub-systems such as, for example, a perception system, a prediction system, a motion planning system, and/or other systems that cooperate to perceive the surrounding environment of the vehicle **105** and determine a motion plan for controlling the motion of the vehicle **105** accordingly. In some implementations, one or more of the perception, prediction, and/or motion planning functions **170A**, **170B**, **170C** can be performed by (and/or combined into) the same system and/or via shared computing resources. In some implementations, one or more of these functions can be performed via difference sub-systems. As further described herein, the autonomy computing system **140** can communicate with the one or more vehicle control systems **150** to operate the vehicle **105** according to the motion plan (e.g., via the vehicle interface **145**, etc.).

(81) The vehicle computing system **110** (e.g., the autonomy computing system **140**) can identify one or more objects that within the surrounding environment of the vehicle **105** based at least in part on the sensor data **135** and/or the map data **160**. The objects perceived within the surrounding environment can be those within the field of view of the sensor(s) **135** and/or predicted to be occluded from the sensor(s) **135**. This can include object(s) not in motion or not predicted to move (static objects) and/or object(s) in motion or predicted to be in motion (dynamic objects/actors). The vehicle computing system **110** (e.g., performing the perception function **170A**, using a perception system, etc.) can process the sensor data **155**, the map data **160**, etc. to obtain perception data **175A**. The vehicle computing system **110** can generate perception data **175A** that is indicative of one or more states or contexts (e.g., current and/or past state(s)) of one or more objects (e.g., actors) that are within a surrounding environment of the vehicle **105**. For example, the perception data **175A** for each object can describe (e.g., for a given time, time period) an estimate of the object's: current and/or past location (also referred to as position); current and/or past speed/velocity; current and/or past acceleration; current and/or past heading; current and/or past orientation; size/footprint (e.g., as represented by a bounding shape, object highlighting, etc.); class (e.g., pedestrian class vs. vehicle class vs. bicycle class, etc.), the uncertainties associated therewith, and/or other state information. The vehicle computing system **110** can utilize one or more algorithms and/or machine-learned model(s) that are configured to identify object(s) based at least in part on the sensor data **155**. This can include, for example, an implicit latent variable model as described herein that is trained to identify object(s) within the surrounding environment of the vehicle **105** and the state data associated therewith. The perception data **175A** can be utilized for the prediction function **170B** of the autonomy computing system **140**. In some implementations, the perception function **170A** and prediction function **170B** may be combined as a joint perception

and prediction task that simultaneously identifies and predicts trajectories of one or more actors.

(82) The vehicle computing system **110** can be configured to predict a motion of the object(s) within the surrounding environment of the vehicle **105**. For instance, the vehicle computing system **110** can generate prediction data **175B** associated with such object(s). The prediction data **175B** can be indicative of one or more predicted future locations of each respective object. For example, the prediction function **170B** can determine a predicted motion trajectory along which a respective object is predicted to travel over time. A predicted motion trajectory can be indicative of a path that the object is predicted to traverse and an associated timing with which the object is predicted to travel along the path. The predicted path can include and/or be made up of a plurality of way points. In some implementations, the prediction data **175B** can be indicative of the speed and/or acceleration at which the respective object is predicted to travel along its associated predicted motion trajectory. The vehicle computing system **110** can utilize one or more algorithms and/or machine-learned model(s) that are configured to predict the future motion of object(s) based at least in part on the sensor data **155**, the perception data **175A**, map data **160**, and/or other data. This can include, for example, an ILVM according to example aspects of the present disclosure that is trained to predict the motion of the object(s) within the surrounding environment of the vehicle **105** based at least in part on the past and/or current state(s) of those objects as well as the environment in which the objects are located (e.g., the lane boundary in which it is travelling, etc.). The prediction data **175B** can be utilized for the motion planning function **170C** of the autonomy computing system **140**.

(83) The vehicle computing system **110** can determine a motion plan for the vehicle **105** based at least in part on the perception data **175A**, the prediction data **175B**, and/or other data. For example, the vehicle computing system **110** can generate motion planning data **175C** indicative of a motion plan. The motion plan can include vehicle actions (e.g., speed(s), acceleration(s), other actions, etc.) with respect to one or more of the objects within the surrounding environment of the vehicle **105** as well as the objects' predicted movements. The motion plan can include one or more vehicle motion trajectories that indicate a path for the vehicle **105** to follow. A vehicle motion trajectory can be of a certain length and/or time range. A vehicle motion trajectory can be defined by one or more way points (with associated coordinates). The planned vehicle motion trajectories can indicate the path the vehicle **105** is to follow as it traverses a route from one location to another. Thus, the vehicle computing system **110** can take into account a route/route data when performing the motion planning function **170C**.

(84) The motion planning function **170C** can implement an optimization algorithm, machine-learned model, etc. that considers cost data associated with a vehicle action as well as other objective functions (e.g., cost functions based on speed limits, traffic lights, etc.), if any, to determine optimized variables that make up the motion plan. The vehicle computing system **110** can determine that the vehicle **105** can perform a certain action (e.g., pass an object, etc.) without increasing the potential risk to the vehicle **105** and/or violating any traffic laws (e.g., speed limits, lane boundaries, signage, etc.). For instance, the vehicle computing system **110** can evaluate the predicted motion trajectories of one or more objects during its cost data analysis to help determine an optimized vehicle trajectory through the surrounding environment. The motion planning function **170C** can generate cost data associated with such trajectories. In some implementations, one or more of the predicted motion trajectories and/or perceived objects may not ultimately change the motion of the vehicle **105** (e.g., due to an overriding factor). In some implementations, the motion plan may define the vehicle's motion such that the vehicle **105** avoids the object(s), reduces speed to give more leeway to one or more of the object(s), proceeds cautiously, performs a stopping action, passes an object, queues behind/in front of an object, etc.

(85) The vehicle computing system **110** can be configured to continuously update the vehicle's motion plan and a corresponding planned vehicle motion trajectories. For example, in some implementations, the vehicle computing system **110** can generate new motion planning data

175C/motion plan(s) for the vehicle **105** (e.g., multiple times per second, etc.). Each new motion plan can describe a motion of the vehicle **105** over the next planning period (e.g., next several seconds, etc.). Moreover, a new motion plan may include a new planned vehicle motion trajectory. Thus, in some implementations, the vehicle computing system **110** can continuously operate to revise or otherwise generate a short-term motion plan based on the currently available data. Once the optimization planner has identified the optimal motion plan (or some other iterative break occurs), the optimal motion plan (and the planned motion trajectory) can be selected and executed by the vehicle **105**.

(86) The vehicle computing system **110** can cause the vehicle **105** to initiate a motion control in accordance with at least a portion of the motion planning data **175C**. A motion control can be an operation, action, etc. that is associated with controlling the motion of the vehicle **105**. For instance, the motion planning data **175C** can be provided to the vehicle control system(s) **150** of the vehicle **105**. The vehicle control system(s) **150** can be associated with a vehicle interface **145** that is configured to implement a motion plan. The vehicle interface **145** can serve as an interface/conduit between the autonomy computing system **140** and the vehicle control systems **150** of the vehicle **105** and any electrical/mechanical controllers associated therewith. The vehicle interface **145** can, for example, translate a motion plan into instructions for the appropriate vehicle control component (e.g., acceleration control, brake control, steering control, etc.). By way of example, the vehicle interface **145** can translate a determined motion plan into instructions to adjust the steering of the vehicle **105** "X" degrees, apply a certain magnitude of braking force, increase/decrease speed, etc. The vehicle interface **145** can help facilitate the responsible vehicle control (e.g., braking control system, steering control system, acceleration control system, etc.) to execute the instructions and implement a motion plan (e.g., by sending control signal(s), making the translated plan available, etc.). This can allow the vehicle **105** to autonomously travel within the vehicle's surrounding environment.

(87) The vehicle computing system **110** can store other types of data. For example, an indication, record, and/or other data indicative of the state of the vehicle (e.g., its location, motion trajectory, health information, etc.), the state of one or more users (e.g., passengers, operators, etc.) of the vehicle, and/or the state of an environment including one or more objects (e.g., the physical dimensions and/or appearance of the one or more objects, locations, predicted motion, etc.) can be stored locally in one or more memory devices of the vehicle **105**. Additionally, the vehicle **105** can communicate data indicative of the state of the vehicle, the state of one or more passengers of the vehicle, and/or the state of an environment to a computing system that is remote from the vehicle **105**, which can store such information in one or more memories remote from the vehicle **105**. Moreover, the vehicle **105** can provide any of the data created and/or store onboard the vehicle **105** to another vehicle.

(88) The vehicle computing system **110** can include the one or more vehicle user devices **180**. For example, the vehicle computing system **110** can include one or more user devices with one or more display devices located onboard the vehicle **105**. A display device (e.g., screen of a tablet, laptop, and/or smartphone) can be viewable by a user of the vehicle **105** that is located in the front of the vehicle **105** (e.g., driver's seat, front passenger seat). Additionally, or alternatively, a display device can be viewable by a user of the vehicle **105** that is located in the rear of the vehicle **105** (e.g., a back passenger seat). The user device(s) associated with the display devices can be any type of user device such as, for example, a table, mobile phone, laptop, etc. The vehicle user device(s) **180** can be configured to function as human-machine interfaces. For example, the vehicle user device(s) **180** can be configured to obtain user input, which can then be utilized by the vehicle computing system **110** and/or another computing system (e.g., a remote computing system, etc.). For example, a user (e.g., a passenger for transportation service, a vehicle operator, etc.) of the vehicle **105** can provide user input to adjust a destination location of the vehicle **105**. The vehicle computing system **110** and/or another computing system can update the destination location of the vehicle **105** and the

route associated therewith to reflect the change indicated by the user input.

(89) The vehicle **105** can be configured to perform vehicle services for one or a plurality of different service entities **185**. A vehicle **105** can perform a vehicle service by, for example and as further described herein, traveling (e.g., traveling autonomously) to a location associated with a requested vehicle service, allowing user(s) and/or item(s) to board or otherwise enter the vehicle **105**, transporting the user(s) and/or item(s), allowing the user(s) and/or item(s) to disembark or otherwise exit the vehicle **105**, etc. In this way, the vehicle **105** can provide the vehicle service(s) for a service entity to a user.

(90) A service entity **185** can be associated with the provision of one or more vehicle services. For example, a service entity can be an individual, a group of individuals, a company (e.g., a business entity, organization, etc.), a group of entities (e.g., affiliated companies), and/or another type of entity that offers and/or coordinates the provision of one or more vehicle services to one or more users. For example, a service entity can offer vehicle service(s) to users via one or more software applications (e.g., that are downloaded onto a user computing device), via a website, and/or via other types of interfaces that allow a user to request a vehicle service. As described herein, the vehicle services can include transportation services (e.g., by which a vehicle transports user(s) from one location to another), delivery services (e.g., by which a vehicle transports/delivers item(s) to a requested destination location), courier services (e.g., by which a vehicle retrieves item(s) from a requested origin location and transports/delivers the item to a requested destination location), and/or other types of services. The vehicle services can be wholly performed by the vehicle **105** (e.g., travelling from the user/item origin to the ultimate destination, etc.) or performed by one or more vehicles and/or modes of transportation (e.g., transferring the user/item at intermediate transfer points, etc.).

(91) An operations computing system **190A** of the service entity **185** can help to coordinate the performance of vehicle services by autonomous vehicles. The operations computing system **190A** can include and/or implement one or more service platforms of the service entity. The operations computing system **190A** can include one or more computing devices. The computing device(s) can include various components for performing various operations and functions. For instance, the computing device(s) can include one or more processors and one or more tangible, non-transitory, computer readable media (e.g., memory devices, etc.). The one or more tangible, non-transitory, computer readable media can store instructions that when executed by the one or more processors cause the operations computing system **190** (e.g., its one or more processors, etc.) to perform operations and functions, such as those described herein matching users and vehicles/vehicle fleets, deploying vehicles, facilitating the provision of vehicle services via autonomous vehicles, etc.

(92) A user **125** can request a vehicle service from a service entity **185**. For example, the user **125** can provide user input to a user device **130** to request a vehicle service (e.g., via a user interface associated with a mobile software application of the service entity **185** running on the user device **130**). The user device **130** can communicate data indicative of a vehicle service request **195** to the operations computing system **190A** associated with the service entity **185** (and/or another associated computing system that can then communicate data to the operations computing system **190A**). The vehicle service request **195** can be associated with a user. The associated user can be the one that submits the vehicle service request (e.g., via an application on the user device **130**). In some implementations, the user may not be the user that submits the vehicle service request. The vehicle service request can be indicative of the user. For example, the vehicle service request can include an identifier associated with the user and/or the user's profile/account with the service entity **185**. The vehicle service request **195** can be generated in a manner that avoids the use of personally identifiable information and/or allows the user to control the types of information included in the vehicle service request **195**. The vehicle service request **195** can also be generated, communicated, stored, etc. in a secure manner to protect information.

(93) The vehicle service request **195** can indicate various types of information. For example, the

vehicle service request **195** can indicate the type of vehicle service that is desired (e.g., a transportation service, a delivery service, a courier service, etc.), one or more locations (e.g., an origin location, a destination location, etc.), timing constraints (e.g., pick-up time, drop-off time, deadlines, etc.), and/or geographic constraints (e.g., to stay within a certain area, etc.). The vehicle service request **195** can indicate a type/size/class of vehicle such as, for example, a sedan, an SUV, luxury vehicle, standard vehicle, etc. The vehicle service request **195** can indicate a product of the service entity **185**. For example, the vehicle service request **195** can indicate that the user is requesting a transportation pool product by which the user would potentially share the vehicle (and costs) with other users/items. In some implementations, the vehicle service request **195** can explicitly request for the vehicle service to be provided by an autonomous vehicle or a human-driven vehicle. In some implementations, the vehicle service request **195** can indicate a number of users that will be riding in the vehicle/utilizing the vehicle service. In some implementations, the vehicle service request **195** can indicate preferences/special accommodations of an associated user (e.g., music preferences, climate preferences, wheelchair accessibility, etc.) and/or other information.

(94) The operations computing system **190A** of the service entity **185** can process the data indicative of the vehicle service request **195** and generate a vehicle service assignment that is associated with the vehicle service request. The operations computing system can identify one or more vehicles that may be able to perform the requested vehicle services to the user **125**. The operations computing system **190A** can identify which modes of transportation are available to a user for the requested vehicle service (e.g., light electric vehicles, human-drive vehicles, autonomous vehicles, aerial vehicle, etc.) and/or the number of transportation modes/legs of a potential itinerary of the user for completing the vehicle service (e.g., single or plurality of modes, single or plurality of legs, etc.). For example, the operations computing system **190A** can determine which autonomous vehicle(s) are online with the service entity **185** (e.g., available for a vehicle service assignment, addressing a vehicle service assignment, etc.) to help identify which autonomous vehicle(s) would be able to provide the vehicle service.

(95) The operations computing system **190A** and/or the vehicle computing system **110** can communicate with one or more other computing systems **190B** that are remote from the vehicle **105**. This can include, for example, computing systems associated with government functions (e.g., emergency services, regulatory bodies, etc.), computing systems associated with vehicle providers other than the service entity, computing systems of other vehicles (e.g., other autonomous vehicles, aerial vehicles, etc.). Communication with the other computing systems **190B** can occur via the network(s) **120**.

(96) FIG. 2 depicts an example service infrastructure **200** according to example embodiments of the present disclosure. The service infrastructure **200** can include one or more systems, interfaces, and/or other components that can be included in an operations computing systems of the service entity for coordinating vehicle services and managing/supporting the autonomous vehicle associated therewith. The service infrastructure **200** can represent, for example, the architecture of a service platform of the operations computing system for coordinating and providing one or more vehicle services (e.g., via autonomous vehicle(s), etc.).

(97) The service infrastructure **200** of an operations computing system can include a first application programming interface platform **205A**, a second application programming interface platform **205B**, and/or a backend system **210** with one or a plurality of backend services **215**. These components can allow the service infrastructure **200** (e.g., the operations computing system) to communicate with one or more autonomous vehicles and/or one or more other systems.

(98) The first application programming interface platform **205A** can facilitate communication with one or more autonomous vehicles of the service entity. For example, as described herein, the service entity may own, lease, etc. a fleet of autonomous vehicles **220A** that can be managed by the service entity (e.g., its backend services) to provide one or more vehicle services. The autonomous

vehicle(s) **220A** can be utilized by the service entity to provide the vehicle service(s) and can be included in the fleet of the service entity. Such autonomous vehicle(s) may be referred to as “service entity autonomous vehicles” or “first party autonomous vehicles.”

(99) The first application programming interface platform **205A** can include a number of components to help facilitate the support, coordination, and management of the first party autonomous vehicles **220A** associated with the service entity. The first application programming interface platform **205A** (e.g., a private platform, etc.) can provide access to one or more backend services **215** that are available to the first party autonomous vehicles **220A**. To help do so, the first application programming interface platform **205A** can include a first API gateway **225A**. The first API gateway **225A** can function as a proxy for application programming interface (API) calls and can help to return an associated response. The first API gateway **225A** can help provide other support functions for the service infrastructure **200** such as, for example, authentication functions, etc.

(100) The first application programming interface platform **205A** can include one or more APIs such as, for example, a first vehicle API **230A**. The first vehicle API **230A** can include a library and/or parameters for facilitating communications between the first party autonomous vehicles **220A** and the backend service(s) **215** of the backend system **210**. For example, the first vehicle API **230A** can be called by a first party autonomous vehicle **220A** and/or another system (e.g., system(s)/platform(s) **250**) to help communicate data, messages, etc. to and/or from an autonomous vehicle and/or another system (e.g., system(s)/platform(s) **250**). The first vehicle API **230A** can provide for communicating such information in a secure, bidirectional manner that allows for expanded processing of data offboard a vehicle, analyzing such data in real time, and/or the like.

(101) The first application programming interface platform **205A** can include first frontend/backend interface(s) **235A**. Each first frontend/backend interface **235A** can be associated with a backend service **215** of the backend system **210**. The first frontend/backend interface(s) **235A** can serve as interface(s) for one client (e.g., an external client such as a first party autonomous vehicle **220A**) to provide data to another client (e.g., a backend service **215**). In this way, the frontend/backend interface(s) **235A** can be external facing edge(s) of the first application programming interface platform **205A** that are responsible for providing secure tunnel(s) for first party autonomous vehicles **220A** (and/or other system(s)/platform(s) **250**) to communicate with the backend system **215** (and vice versa) so that a particular backend service can be accessed by a particular first party autonomous vehicle **220A** (and/or other system(s)/platform(s) **250**).

(102) In some implementations, the first application programming interface platform **205A** can include one or more first adapters **240A**, for example, to provide compatibility between one or more first frontend/backend interfaces **235A** and one or more of the API(s) associated with the first application programming interface platform **205A** (e.g., vehicle API **230A**). The first adapter(s) **240A** can provide upstream and/or downstream separation between particular infrastructure components, provide or assist with data curation, flow normalization and/or consolidation, etc.

(103) The second application programming interface platform **205B** (e.g., a public platform, etc.) can facilitate communication with one or more autonomous vehicles of a third party vehicle provider. As described herein, a third party vehicle provider can be an entity that makes one or more of its autonomous vehicles available to the service entity for the provision of vehicle services. This can include, for example, an individual, an original equipment manufacturer (OEM), a third party vendor, or another entity that places its autonomous vehicle(s) online with the service platform of the service entity such that the autonomous vehicle(s) can provide vehicle services of the service entity. These autonomous vehicles may be referred to as “third party autonomous vehicles” and are shown in FIG. 2 as third party autonomous vehicles **220B**. Even though such autonomous vehicles may not be included in the fleet of autonomous vehicles of the service entity, the service infrastructure **200** (e.g., of the service entity's service platform, etc.) can allow the third party autonomous vehicles **220B** to provide vehicle services offered by the service entity, access

the one or more backend services **215** of the backend system **210**, etc.

(104) The second application programming interface platform **205B** can allow the service platform to communicate directly or indirectly with autonomous vehicle(s). In some implementations, a third party autonomous vehicle **220B** may call an API of, send data/message(s) to, receive data/message(s) from/directly through, etc. the second application programming interface platform **205B**.

(105) Additionally, or alternatively, another computing system can serve as an intermediary between the third party autonomous vehicles **220B** and the second application programming interface platform **205B** (and the service platform associated therewith). For example, the service infrastructure **200** can be associated with and/or in communication with one or more third party vehicle provider computing systems, such as a vehicle provider X computing system **245A** and a vehicle provider Y computing system **245B**. Each third party vehicle provider X, Y can have its own, separate third party autonomous fleet including respective third party autonomous vehicles **220B**. The third party vehicle provider computing systems **245A**, **245B** can be distinct and remote from the service infrastructure **200** and provide for management of vehicles associated with that particular third party vehicle provider. As shown in FIG. 2, a third party vehicle provider computing system **245A**, **245B** can include its own backends and/or frontends for communicating with other systems (e.g., third party autonomous vehicle(s) **220B**, operations computing system, etc.).

(106) The third party computing system **245A**, **245B** associated with a particular third party autonomous vehicle fleet can serve as the communication intermediary for that fleet. For example, third party autonomous vehicles **220B** associated with third party vehicle provider X can communicate with the third party vehicle provider X computing system **245A** which can then communicate with the service infrastructure **200** (e.g., to access the available backend services **215**) via the second application programming interface platform **205B**. Data from the service infrastructure **200** (e.g., the backend services **215**) can be communicated to the vehicle provider X computing system **245A** (e.g., via the second application programming interface platform **235B**) and then to the third party autonomous vehicles **220B** associated with third party vehicle provider X. In another example, third party autonomous vehicles **220B** associated with third party vehicle provider Y can communicate with the third party vehicle provider Y computing system **245B** which can then communicate with the service infrastructure **200** (e.g., to access the available backend services **215**) via the second application programming interface platform **205B**. Data from the service infrastructure **200** (e.g., the backend services **215**) can be communicated to the third party vehicle provider Y computing system **245B** (e.g., via the second application programming interface platform **205B**) and then to the third party autonomous vehicles **220B** associated with third party vehicle provider Y.

(107) The second application programming interface platform **205B** can include a number of components to help facilitate the support, coordination, and management of the third party autonomous vehicles **220B** associated with the third party vehicle providers. The second application programming interface platform **205B** can provide access to one or more backend services **215** that are available to the third party autonomous vehicles **220B**. To help do so, the second application programming interface platform **205B** can include a second API gateway **225B**. The second API gateway **225B** can function as a proxy for application programming interface (API) calls and can help to return an associated response. The second API gateway **225B** can help provide other support functions for the service infrastructure **200** such as, for example, authentication functions, etc.

(108) The second application programming interface platform **205B** can include one or more APIs such as, for example, a second vehicle API **230B**. The second vehicle API **230B** can include a library and/or parameters for facilitating communications between the third party autonomous vehicles **220B** and the backend service(s) **215** of the backend system **210**. For example, the second

vehicle API **230B** can be called by a third party autonomous vehicle **220B** and/or another system (e.g., a third party vehicle provider computing system **245B**, etc.) to help communicate data, messages, etc. to and/or from an autonomous vehicle. The second vehicle API **230B** can provide for communicating such information in a secure, bidirectional manner.

(109) The second application programming interface platform **205B** can include second frontend/backend interface(s) **235B**. Each of the second frontend/backend interface(s) **235B** can be associated with a backend service **215** of the backend system **210**. The second frontend/backend interface(s) **235B** can serve as interface(s) for one client (e.g., an external client such as a third party autonomous vehicle **220B**, a third party vehicle provider computing system **245A**, **245B**) to provide data to another client (e.g., a backend service **215**). In this way, the second frontend/backend interface(s) **235B** can be external facing edge(s) of the second application programming interface platform **205B** that are responsible for providing secure tunnel(s) for third party autonomous vehicles **220B** (and/or other intermediary systems) to communicate with the backend system **210** (and vice versa) so that a particular backend service **215** can be utilized. In some implementations, the second application programming interface platform **205B** can include one or more second adapters **240B**, for example, to provide compatibility between one or more second frontend/backend interfaces **235B** and one or more of the API(s) associated with the second application programming interface platform **205B** (e.g., vehicle API **230B**).

(110) In some implementations, the first party autonomous vehicles **220A** can utilize the second application programming interface platform **205B** to access/communicate with the service platform/backend service(s) **215**. This can allow for greater accessibility and/or back-up communication options for the first party autonomous vehicles **220A**.

(111) The backend system **210** can host, store, execute, etc. one or more backend services **215**. The backend service(s) **215** can be implemented by system client(s), which can include hardware and/or software that is remote from the autonomous vehicles and that provide a particular service to an autonomous vehicle. The backend service(s) **215** can include a variety of services that help coordinate the provision of vehicle service(s) and support the autonomous vehicles and/or the third party vehicle providers performing/providing those vehicle service(s).

(112) For example, the backend service(s) **215** can include a matching service that is configured to match an autonomous vehicle and/or an autonomous vehicle fleet with a service request for vehicle services. Based on a match, the matching service can generate and communicate data indicative of a candidate vehicle service assignment (indicative of the requested vehicle service) for one or more autonomous vehicles. In some implementations (e.g., for first party autonomous vehicle(s) **220A**), the candidate vehicle service assignment can include a command that a first party autonomous vehicle **220A** is required to accept, unless it would be unable to safely or fully perform the vehicle service. In some implementations (e.g., for third party autonomous vehicle(s) **220B**), the candidate vehicle service assignment can include a request or offer for one or more autonomous vehicles to provide the vehicle service. The candidate vehicle service assignment can be communicated to one or more third party vehicle provider computing systems **245A**, **245B** and/or one or more autonomous vehicle(s) **220B** (e.g., via the second application programming interface platform **205B**) and/or one or more autonomous vehicle(s) **220A** (e.g., via the first application programming interface platform **205A**). The candidate vehicle service assignment can be accepted or rejected. If accepted, an autonomous vehicle **220A**, **220B** can be associated (e.g., assigned to service, etc.) with the vehicle service assignment. The vehicle service assignment can include data indicative of the user, a route, an origin location for the vehicle service, a destination location for the vehicle service, service parameters (e.g., time restraints, user accommodations/preferences, etc.), and/or any other information associated with a vehicle service.

(113) The backend service(s) **215** can include an itinerary service. The itinerary service can be configured to maintain, update, track, etc. a data structure indicative of one or more task(s) and/or candidate task(s) associated with (and/or potentially associated with) a particular autonomous

vehicle, autonomous vehicle fleet, and/or vehicle provider. The tasks can include, for example, vehicle service assignments for providing vehicle services and/or tasks associated with an activity other than the performance of a vehicle service. For example, the tasks can include: a testing task (e.g., for testing and validating autonomy software, hardware, etc.); a data acquisition task (e.g., acquiring sensor data associated with certain travel ways, etc.); a re-positioning task (e.g., for moving an idle vehicle between vehicle service assignments, to high demand areas, etc.); a circling task (e.g., for travelling within the current geographic area in which a vehicle is located (e.g., circle the block or neighborhood), etc.); a maintenance task (e.g., for instructing travel to a service depot to receive maintenance, etc.); a re-fueling task; a vehicle assistance task (e.g., where a vehicle travels to assist another vehicle, etc.); a deactivation task (e.g. going offline such that a vehicle, fleet of vehicles, or vehicle providers no longer accept service request, etc.); a parking task; and/or other types of tasks. The itinerary service can maintain an itinerary for an autonomous vehicle, fleet, vehicle provider, etc. The itinerary can serve as a queue for the various tasks. In some implementations, the tasks can be associated with a priority or order for which they are deployed to an autonomous vehicle, fleet, vehicle provider, etc.

(114) In some implementations, the vehicle service assignment can be associated with a multi-modal vehicle service. For example, the user may request and/or be provided a multi-modal user itinerary by which the user is to travel to the user's ultimate destination via two or more types of transportation modalities (e.g., ground based vehicle, aerial vehicle, public transit, etc.). As such, the origin location and/or destination location identified in the vehicle service assignment may include intermediate locations (e.g., transfer points) along the user's multi-modal itinerary.

(115) The backend service(s) **215** can include a deployment service that communicates tasks for an autonomous vehicle to complete. For example, the deployment service can communicate data indicative of a vehicle service assignment and/or another task to an autonomous vehicle (or an intermediary system). The deployment service can communicate such data to an autonomous vehicle (or an intermediary system) based at least in part on the itinerary associated therewith. By way of example, the highest priority task and/or the task that is next in order can be deployed.

(116) The backend service(s) **215** can include a routing service. The routing service can be configured to provide an autonomous vehicle with a route for a vehicle service and/or another task. The route can be based at least in part on factors associated with the geographic area in which the autonomous vehicle is (or will be) travelling (e.g., roadways, weather, traffic, events, etc.). Additionally, or alternatively, the route can be based at least in part the autonomy capabilities of the autonomous vehicle (e.g., ability to complete an unprotected left-hand turn, U-turn, etc.). In some implementations, the routing service can be configured to assign, coordinate, monitor, adjust, etc. one or more designated pick-up and/or drop-off zones for the vehicle service(s). The routing service can be available to first party autonomous vehicles **220A**. In addition, or alternatively, the routing service can be available to third party autonomous vehicles **220B** if permitted/requested by an associated third party vehicle provider.

(117) The backend service(s) **215** can include a rider experience service. The rider experience service can be configured to communicate data to a rider associated with the vehicle service. This can include, for example, upcoming vehicle actions, routes, drop-off zones, user adjustable vehicle conditions (e.g., music, temperature, etc.). Such information can be presented via a display device of an onboard tablet associated with an autonomous vehicle, a user device associated with the rider, etc. through a software application associated with the service entity.

(118) The backend service(s) **215** can include a remote assistance service. The remote assistance service can be configured to provide remote assistance to an autonomous vehicle and/or a user (e.g., a rider associated with the vehicle service, etc.). For example, a remote assistance operator can take over control of one or more vehicle operations and/or otherwise assist an autonomous vehicle during the one or more vehicle operations. By way of example, a remote assistance operator can remotely control the navigation of an autonomous vehicle to navigate the vehicle

around/past an unexpected obstruction in a travel way (e.g., a fallen tree, etc.). In another example, the remote assistance operator can communicate with a user (e.g., via the onboard tablet, user's phone, etc.) in the event that the user is in need of help.

(119) The backend service(s) **215** can include a simulation/testing service. The simulation/testing service can help facilitate vehicle provider integration with the service platform. For example, simulation/testing service can provide testing environments for vehicle providers to simulate communications and/or the performance of vehicle services using the service infrastructure **200**.

(120) The backend service(s) **215** can include one or more other services. This can include, for example, payment services, vehicle rating services, health and maintenance services, software update/deployment services, and/or other services.

(121) In some implementations, one or more backend services **215** that are available to the first party autonomous vehicles **220A** (e.g., via the first application programming interface platform **205A**) may not be available to the third party autonomous vehicles **220B** (e.g., via the second application programming interface platform **205B**), and vice versa. For example, a software update/deployment service for the first party autonomous vehicles **220A** may not be accessible or suitable for a third party autonomous vehicle **220B** that utilizes the onboard autonomy software of a third party vehicle provider (not the service entity). As such, a software update/deployment backend service may not be able to communicate with a third party autonomous vehicle **220B** and/or vice versa.

(122) In some implementations, the service infrastructure **200** can include a test platform for validating and vetting end-to-end platform functionality, without use of a real vehicle on the ground. For example, the test platform can simulate trips with human drivers and/or support fully simulated trip assignment and/or trip workflow capabilities. For example, the test platform can simulate and monitor data traffic through the service infrastructure **200** to ensure proper functioning. In some implementations, the testing platform can access the simulation/testing backend to help facilitate a test or simulation.

(123) In some implementations, the service infrastructure **200** can utilize a plurality of software development kits (SDKs) that help provide access to the first and second application programming interface platforms **205A**, **205B**. All (or a portion of) external communication with the platforms can be done via the SDKs. For example, the SDKs can include a first SDK (e.g., private SDK) and a second SDK (e.g., public SDK) and specific endpoints to facilitate communication with the first and second application programming interface platforms **205A**, **205B**, respectively. In some implementations, the first party autonomous vehicle(s) **220A** (and/or a test platform) can use both the first and second SDKs, whereas the third party autonomous vehicles **220B** and/or the third party vehicle provider computing systems **245A**, **245B** can use only the second SDK and associated endpoints. In some implementations, the SDKs can provide a single entry point, which can improve consistency across both the service provider fleet and the third party entity fleet(s). As an example, a second SDK can provide secured access to the second application interface platform **205B** and access to capabilities such as vehicle service assignments, routing, and/or the like. The first SDK can be accessed by the first party autonomous vehicles **220A** and provide access to capabilities including those available only to the first party autonomous vehicles **220A**.

(124) In some implementations, the SDKs can include a command-line interface to provide an entry point into the SDK components and act as a gateway for SDK related work, integration, testing, and authentication. For example, the command-line tools can provide for bootstrapping, managing authentication, updating SDK version, testing, debugging, and/or the like. In some implementations, a command-line interface can require an authentication certificate before being able to bootstrap an SDK, download components, and/or access a service entity's services. For example, based on the authentication certificate, a command-line interface can determine which version of the SDK to which to provide access. In some implementations, SDKs can be implemented onboard a first or third party autonomous vehicle **220A**, **220B** and/or a third party

vehicle provider computing system **245A**, **245B**.

(125) In some implementations, the service infrastructure **200** can facilitate communication between the service platform and one or more other system(s)/platform(s) **250** associated with the service entity/operations computing system. By way of example, the service entity may have (e.g., the operations computing system may include, etc.) one or more other system(s)/platform(s) **250** that can help indicate what services/vehicles are available to a user or other system, coordinate the provision of vehicle services by human-driven vehicles, and/or are specifically associated with certain types of services (e.g., delivery services, aerial transport services, etc.). The other system(s)/platform(s) **250** may communicate with the service platform utilizing the service infrastructure **200** (e.g., first application programming interface platform **205A**, second application programming interface platform **205B**, etc.) to determine, for example, whether any autonomous vehicles would be available to the user for any potential vehicle services.

(126) FIG. 3 depicts a flow diagram **300** of an example feature extraction according to example aspects of the present disclosure. The feature extraction can be performed using sensor data **302** and/or map data **304**. In some implementations, sensor data **302** can correspond to sensor data **155** of FIG. 1, while map data **304** can correspond to map data **160** of FIG. 1. For instance, one or more sensor(s) (e.g., onboard an autonomous vehicle) can be configured to acquire sensor data **302** (e.g., one or more sensor observations). The sensor(s) can be external sensors configured to acquire external sensor data and/or internal sensors configured to acquire internal sensor data. For instance, a surrounding environment of a vehicle can include/be represented in the field of view of the sensor(s). For instance, the sensor data **302** can be or can include image and/or other data of the environment outside of the vehicle and within a range and/or field of view of one or more of the sensor(s). For instance, the sensor data **302** can be or can include one or more of Light Detection and Ranging (LIDAR) data, one Radio Detection and Ranging (RADAR) data, image data from one or more cameras (e.g., visible spectrum cameras, infrared cameras, etc.), detection and/or ranging data from one or more motion sensors, audio data from one or more audio sensors (e.g., microphones, etc.), and/or other sensor data from any suitable types of imaging capture devices and/or sensors.

(127) The map data **304** can provide detailed information about the surrounding environment of an autonomous vehicle and/or the geographic area in which an autonomous vehicle was, is, and/or will be located. For example, the map data **304** can provide information regarding: the identity and location of different roadways, road segments, buildings, or other items or objects (e.g., lampposts, crosswalks and/or curb); the location and directions of traffic lanes (e.g., the location and direction of a parking lane, a turning lane, a bicycle lane, or other lanes within a particular roadway or other travel way and/or one or more boundary markings associated therewith); traffic control data (e.g., the location and instructions of signage, traffic lights, and/or other traffic control devices); obstruction information (e.g., temporary or permanent blockages, etc.); event data (e.g., road closures/traffic rule alterations due to parades, concerts, sporting events, etc.); nominal vehicle path data (e.g., indicate of an ideal vehicle path such as along the center of a certain lane, etc.); and/or any other map data that provides information that assists a vehicle computing system in processing, analyzing, and perceiving its surrounding environment and its relationship thereto. In some implementations, the map data **304** can include high definition map data. In some implementations, the map data **304** can include sparse map data indicative of a limited number of environmental features (e.g., lane boundaries, etc.). In some implementations, the map data can be limited to geographic area(s) and/or operating domains in which an autonomous vehicle (or autonomous vehicles generally) may travel (e.g., due to legal/regulatory constraints, autonomy capabilities, and/or other factors).

(128) The sensor data **302** and/or the map data **304** (e.g., one or more scene observations) can be provided to a scene feature extraction model **306**. For instance, the scene feature extraction model **306** can be or can include a backbone network, such as a convolutional backbone network. For

instance, the backbone network can be a lightweight backbone network adapted for feature extraction. In some implementations, two separate streams can be instantiated such that the voxelized LiDAR and rasterized map are processed separately. The resulting features from both streams can then be concatenated feature-wise (e.g., if they share the same spatial resolution) and fused by a convolutional header. These extracted features can inform both the downstream detection and motion forecasting networks. In some implementations, the scene feature extraction model **306** (e.g., the backbone network) includes two convolutional layers to output a classification or confidence score and/or a bounding box for each anchor location (e.g., each scene feature). These features are eventually reduced to the final set of candidates by applying non-maximal suppression (NMS) and finally thresholding low probability detections (given by the desired common recall). In some implementations, a backbone network along with features for object detection and per actor feature extraction are provided. The scene feature extraction model **306** can produce scene features **308** (also referred to as global features).

(129) The one or more scene features **308** can be provided to an actor feature recognition model. As used herein, unless indicated otherwise, an “actor feature recognition model” can refer to any collection of one or more operations to recognize one or more actor features corresponding to actors or traffic participants from a broader one or more scene features. For instance, the actor feature recognition model can parse the scene features into actor features that correspond to an actor. For instance, the one or more actor features can include data descriptive of an actor context of one or more traffic participants. For instance, rotated region of interest (ROI) align can be applied to the scene features **308** to extract (e.g., fixed size) spatial feature maps **309** for bounding boxes with arbitrary shapes and rotations from the scene features **308**. The region around each spatial feature map **309** can be pooled to produce pooled actor features **310**. For instance, a region around each actor in its frame can be pooled, such as with an axis defined by the actor's centroid orientation. The pooled actor features **310** may be larger than the eventual actor features. The pooled actor features **310** can be downsampled by applying one or more downsampling convolutional neural networks (CNNs) **312**. As one example, a 4-layer down-sampling convolutional network can be applied. Additionally, and/or alternatively, the pooled actor features **310** can be max-pooled along spatial dimensions to reduce each pooled actor feature **310** to a respective actor feature **314** (e.g., an actor context). For instance, after the downsampling CNN **312** is applied, it can be followed by max-pooling along the spatial dimensions to reduce the pooled actor features **310** to a desired dimensional feature vector **314** per actor. One example convolutional network **312** uses a dilation factor of 2 for the convolutional layers to enlarge the receptive field for the per-actor features, which can improve performance.

(130) FIG. 4 depicts a diagram **400** of an example latent distribution according to example aspects of the present disclosure. For instance, a latent distribution **404** can be established to map from a plurality of actor contexts **402** (e.g., $x_1, x_2 \dots$) to a plurality of predicted future trajectories **406** (e.g., $y_1, y_2 \dots$). As one example, a latent prior model (e.g., latent prior model **510** of FIG. 5) can be configured to map the actor contexts **402** to the latent distribution **404**. As illustrated in FIG. 4, the latent distribution **404** can include a plurality of latent variables **405** (e.g., $z_1, z_2, \dots$). For instance, in some implementations, each latent variable **404** can correspond to a unique actor represented by a unique actor context **402** and/or a unique future trajectory **406**. For instance, in some implementations, the latent distribution **404** can be partitioned as a distributed representation where each latent variable **405** in the scene latent distribution **404** is anchored to a respective actor, such as anchored to a particular node in an interaction graph with traffic participants as nodes. The distributed representation may provide the benefit of naturally scaling the capacity of the latent distribution as the number of actors grows. Furthermore, the anchoring may give the model an inductive bias that eases the learning of a scene latent distribution. Intuitively, each anchored latent variable **405** encodes unobserved dynamics most relevant to its respective actor, including interactions with neighboring actors and traffic rules that apply in its locality. For instance, each

latent variable **405** can be represented as a diagonal multivariate Gaussian. Each latent variable **405** can be conditioned on all actors such that the latent distribution is not marginally independent across actors, although factorized. For instance, despite anchoring each partition of the scene latent to an actor, each individual latent variable **405** can contain information about the full scene, since each final node representation is dependent on the entire input because of the message propagation in the fully-connected interaction graph. In some implementations, scene interaction modules including one or more graph neural networks can be used to produce the latent distribution **404** based at least in part on the actor contexts **402** and/or to produce predicted trajectories **406** based at least in part on samples of the latent distribution **404**. For instance, in some implementations, a decoder model (e.g., decoder model **516** of FIG. 5) can map samples of the latent distribution **404** to the predicted trajectories **406**.

(131) FIG. 5 depicts a flow diagram of an example implicit latent variable model **500** according to example aspects of the present disclosure. The implicit latent variable model **500** can receive as input sensor data **502** and/or map data **504**. In some implementations, sensor data **502** can correspond to sensor data **155** of FIG. 1 and/or sensor data **302** of FIG. 3, while map data **504** can correspond to map data **160** of FIG. 1 and/or map data **304** of FIG. 3. For instance, one or more sensor(s) (e.g., onboard an autonomous vehicle) can be configured to acquire sensor data **502** (e.g., one or more sensor observations). The sensor(s) can be external sensors configured to acquire external sensor data and/or internal sensors configured to acquire internal sensor data. For instance, a surrounding environment of a vehicle can include/be represented in the field of view of the sensor(s). For instance, the sensor data **502** can be or can include image and/or other data of the environment outside of the vehicle and within a range and/or field of view of one or more of the sensor(s). For instance, the sensor data **502** can be or can include one or more of Light Detection and Ranging (LIDAR) data, one Radio Detection and Ranging (RADAR) data, image data from one or more cameras (e.g., visible spectrum cameras, infrared cameras, etc.), detection and/or ranging data from one or more motion sensors, audio data from one or more audio sensors (e.g., microphones, etc.), and/or other sensor data from any suitable types of imaging capture devices and/or sensors.

(132) The map data **504** can provide detailed information about the surrounding environment of an autonomous vehicle and/or the geographic area in which an autonomous vehicle was, is, and/or will be located. For example, the map data **504** can provide information regarding: the identity and location of different roadways, road segments, buildings, or other items or objects (e.g., lampposts, crosswalks and/or curb); the location and directions of traffic lanes (e.g., the location and direction of a parking lane, a turning lane, a bicycle lane, or other lanes within a particular roadway or other travel way and/or one or more boundary markings associated therewith); traffic control data (e.g., the location and instructions of signage, traffic lights, and/or other traffic control devices); obstruction information (e.g., temporary or permanent blockages, etc.); event data (e.g., road closures/traffic rule alterations due to parades, concerts, sporting events, etc.); nominal vehicle path data (e.g., indicate of an ideal vehicle path such as along the center of a certain lane, etc.); and/or any other map data that provides information that assists a vehicle computing system in processing, analyzing, and perceiving its surrounding environment and its relationship thereto. In some implementations, the map data **504** can include high definition map data. In some implementations, the map data **504** can include sparse map data indicative of a limited number of environmental features (e.g., lane boundaries, etc.). In some implementations, the map data can be limited to geographic area(s) and/or operating domains in which an autonomous vehicle (or autonomous vehicles generally) may travel (e.g., due to legal/regulatory constraints, autonomy capabilities, and/or other factors).

(133) The sensor data **502** and/or the map data **504** (e.g., one or more scene observations) can be provided to a scene feature extraction model **506**. For instance, the scene feature extraction model **506** can be or can include a backbone network, such as a convolutional backbone network. For

instance, the backbone network can be a lightweight backbone network adapted for feature extraction. In some implementations, two separate streams can be instantiated such that the voxelized LiDAR and rasterized map are processed separately. The resulting features from both streams can then be concatenated feature-wise (e.g., if they share the same spatial resolution) and fused by a convolutional header. These extracted features can inform both the downstream detection and motion forecasting networks. In some implementations, the scene feature extraction model **506** (e.g., the backbone network) includes two convolutional layers to output a classification or confidence score and/or a bounding box for each anchor location (e.g., each scene feature). These features are eventually reduced to the final set of candidates by applying non-maximal suppression (NMS) and finally thresholding low probability detections (given by the desired common recall). In some implementations, a backbone network along with features for object detection and per actor feature extraction are provided. The scene feature extraction model **506** can produce scene features (also referred to as global features).

(134) The scene features can then be provided to an actor feature recognition model **508**. The actor feature recognition model **508** can parse the scene features into actor features that correspond to an actor. For instance, the one or more actor features can include data descriptive of an actor context of one or more traffic participants. For instance, the actor feature recognition model **508** can be configured to extract spatial feature maps for bounding boxes from the one or more scene features by rotated region of interest (ROI) align. Rotated ROI align can be applied to extract (e.g., fixed size) spatial feature maps for bounding boxes with arbitrary shapes and rotations from the scene features (e.g., the global feature map extracted by the backbone). For instance, rotated ROI align can provide actor contexts for each actor.

(135) Additionally and/or alternatively, the actor feature recognition model **508** can be configured to pool a region around each spatial feature map to produce pooled actor features. For instance, a region around each actor in its frame can be pooled, such as with an axis defined by the actor's centroid orientation. The pooled actor features may be larger than the eventual actor features.

(136) Additionally and/or alternatively, the actor feature recognition model **508** can be configured to downsample the pooled actor features by applying one or more downsampling convolutional neural networks. As one example, a 4-layer down-sampling convolutional network can be applied.

(137) Additionally and/or alternatively, the actor feature recognition model **508** can be configured to max-pool along spatial dimensions to reduce each pooled actor feature to a respective actor feature of the one or more actor features. For instance, after the downsampling CNN is applied, it can be followed by max-pooling along the spatial dimensions to reduce the feature map to a desired dimensional feature vector per actor. One example convolutional network uses a dilation factor of 2 for the convolutional layers to enlarge the receptive field for the per-actor features, which can improve performance. One or more actor features can be obtained from the actor feature recognition model **508**.

(138) The scene data including the one or more actor features can then be provided to a latent prior model **510**. The latent prior model **510** can be configured to generate scene latent data in response to receipt of scene data. In some implementations, the latent prior model **510** can be or can include a scene interaction module including one or more graph neural networks. The scene latent data **512** can be or can include one or more latent variables. In some implementations, the scene latent data **512** can include a latent distribution that is partitioned into one or more latent variables. As one example, the one or more latent variables of the scene latent data **512** can be partitioned with respect to actor contexts, such as depicted in FIG. 4 by latent variables **405**. For instance, in some implementations, the one or more latent variables can be respective to the one or more actor features such that each actor feature has an associated latent variable of the scene latent data **512** that is anchored to the actor feature. In some implementations, the one or more latent variables can be or can include one or more continuous latent variables. Scene latent data **512** can be obtained from the latent prior model **510**.

(139) Latent sample data **514** can be sampled from the scene latent data **512**. For instance, the scene latent data **512** (e.g., the one or more latent variables) can define a latent (e.g., a latent distribution) that can be sampled to produce latent samples **514** of the scene latent data **512**. The latent sample data **514** can define a possible future for the actors (e.g., the actor features).

(140) The latent sample data **514** can be provided to a decoder model **516**. The decoder model **516** can be configured to decode the latent sample data **514** into a motion forecast **518** including one or more predicted trajectories of the one or more actor features. For instance, the decoder model **516** can produce a motion forecast **518** from latent samples. The motion forecast **518** including one or more predicted trajectories of one or more actor features can be received from the decoder model **516**.

(141) In some implementations, the decoder model **516** can be or can include a deterministic decoder model **516**. For instance, the decoder can be or can include a deterministic mapping to implicitly characterize the joint probability, such as opposed to explicitly representing it in a parametric form. This approach can provide for evaluation without factorizing the joint distribution and thus can sidestep potential associated challenges. In particular, the deterministic decoder can be highly beneficial for generating socially consistent trajectories. In this framework, generating scene-consistent future trajectories across actors is simple and highly efficient, as it may be performed with only one stage of parallel sampling.

(142) In some implementations, the decoder model **516** can be or can include a scene interaction module including one or more graph neural networks. For instance, the decoder including a scene interaction module can predict a realization of the future at the scene level via message parsing. As an example, each actor trajectory can be established with respect to samples from each latent variable of the scene latent data **512** and/or each actor feature. This can provide for improved reasoning about multi-agent interactions such as car following, yielding, etc. For instance, each actor context can be initialized as a node in the decoder SIM. After a round of message parsing, each node can then contain an updated representation of a respective actor that takes into account the underlying dynamics of the scene summarized in the latent distribution. Finally, the trajectory sample for the actor can be deterministically decoded, such as by the output function of the SIM. For instance, the output function can be the deterministic mapping. This can be performed without requiring any additional sampling steps. The trajectory-level scene sample can thus be a collection of all actor trajectories.

(143) In some implementations, the decoder model **516** can be or can include a specified and tractable conditional likelihood. In this implementation, many tools are available for inference and learning. As one example, variational inference, such as the variational auto-encoder (VAE), can be used.

(144) In some implementations, the decoder model **516** can be characterized via a stochastic sampling procedure where a conditional dependence is not specified. In this case, likelihood-free inference methods can be employed for learning. As examples, density estimation can be implemented by comparison using either density ratio (GAN) or density difference (MMD).

(145) FIG. 6 depicts a flow diagram of training an example implicit latent variable model **600** according to example aspects of the present disclosure. For instance, FIG. 6 illustrates components of a training loss **630** that can be established during training of the model **600**. As illustrated in FIG. 6, sensor data **602** and/or map data **604** can be provided to a feature extraction model **606** to produce one or more actor features **608**, as described herein. The one or more actor features **608** can be provided to an object detector **609** to detect objects (e.g., actors). The sensor data **602** and/or map data **604** can be at least a portion of a training dataset including ground truth data, such as ground truth data including a ground truth context of the one or more actor features **608**. The ground truth context can be descriptive of a known context of the actor features **608**, such as a known position, pose, velocity, etc. For instance, the ground truth data can include training features including known contexts of the actor features. A feature loss can be determined between one or

more predicted features from the object detector **609** and the training features. For instance, the feature loss can include a cross-entropy loss **632** between the one or more predicted features and one or more training features of the ground truth data. Additionally and/or alternatively, the feature loss can include a regression loss **634** between bounding boxes of the one or more predicted features and bounding boxes of the ground truth data. For instance, in some implementations, the feature loss includes a cross entropy loss **632** applied to detection scores of the predicted features and/or a smooth L1 **634** applied to the bounding box regression targets of the one or more predicted features. The training loss **630** can be or can include the cross entropy loss **632** and/or the regression loss **634**.

(146) The training examples labeled with ground truth data can be provided to a latent encoder model **610**. The latent encoder model **610** can be configured to produce a first latent distribution **612** in response to receipt of the one or more training examples and the ground truth data. For instance, in some cases, integration over the scene latent distribution is intractable, in which case amortized variational inference can be used. For instance, by introducing an encoder distribution to approximate the true posterior, the learning problem can be reformulated as a maximization of the Evidence Lower BOund (ELBO). In some implementations, the latent encoder model **610** can include a scene interaction module. For instance, after running one round of message passing, the scene interaction module can predict the distribution over latent variables.

(147) The encoder model **610** can approximate the true posterior latent distribution. This model may also be called a recognition network. Because the encoder model **610** receives the ground truth data (e.g., the target output) as an input, it can recognize scene dynamics that are unobserved by the latent prior model **614**. In this manner, the encoder may only be used during training, since it requires access to the ground-truth future trajectories. For instance, the encoder may be omitted from deployed models and/or included at an online training system.

(148) Additionally, the one or more training examples can be provided to a latent prior model **614**. The latent prior model **614** can be configured to produce a second latent distribution **616** in response to receipt of the one or more training examples. For instance, the latent prior model **614** can be agnostic to the ground truth data such that the latent prior model **614** is usable during inference (e.g., during prediction after ground truth data has already been used to train the prior model **614**). The latent prior model **614** may be utilized during inference.

(149) A divergence loss **636** can be determined based at least in part on the first latent distribution **612** and the second latent distribution **616**. For instance, in some implementations, the divergence loss **636** can be or can include a KL divergence loss **636** between the first latent distribution **612** and the second latent distribution **616**. The training loss **630** can be or can include the divergence loss **636**.

(150) Samples from the first latent distribution **612** and/or the second latent distribution **616** can be provided to decoder model **618** to produce motion forecast(s) **620**. The motion forecast(s) **620** can be used to produce a reconstruction loss **638**. For instance, the reconstruction loss **638** can be obtained by sampling (e.g., by a computing system) the first latent distribution to produce one or more first latent samples. The first latent distribution can include first sample data, such as a first set of samples from the latent distribution. Determining the reconstruction loss **638** can further include sampling (e.g., by the computing system) the second latent distribution to produce one or more second latent samples. Determining the reconstruction loss **638** can further include providing (e.g., by the computing system) the one or more first latent samples to a decoder model **618** and receiving (e.g., by the computing system) one or more first predicted trajectories **620** from the decoder model **618**. The first predicted trajectories **620** can be representative of a first possible future of the actor features. Determining the reconstruction loss **638** can further include providing (e.g., by the computing system) the one or more second latent samples to the decoder model **618** and receiving (e.g., by the computing system) one or more second predicted trajectories **620** from the decoder model **618**. The second predicted trajectories **620** can be representative of a second

possible future of the actor features that desirably matches the first possible future. Determining the reconstruction loss **638** can include determining (e.g., by the computing system) a reconstruction loss **638** based at least in part on the one or more first predicted trajectories **620** and the one or more second predicted trajectories **620**. For instance, the reconstruction loss **638** can be based at least in part on a difference between the one or more first predicted trajectories **620** and the one or more second predicted trajectories **620**. In some implementations, the reconstruction loss **638** can include a Huber loss. The training loss **630** can be or can include the reconstruction loss **638**.

(151) FIG. 7 depicts a flow chart diagram of an example method **700** for determining scene-consistent motion forecasts from sensor data according to example aspects of the present disclosure. One or more portions of the method **700** described herein can be implemented as an algorithm on the hardware components of the devices described herein (e.g., as in FIGS. 1, 2, 12, 13) to, for example, generate scenarios using parametric modeling and conducting simulations to test autonomous vehicle (e.g., autonomous vehicle **105** in FIG. 1) using those generated scenarios. Although FIG. 7 depicts steps performed in a particular order for purposes of illustration and discussion, method **700** of FIG. 7 is not limited to the particularly illustrated order or arrangement. The various steps of the methods disclosed herein can be omitted, rearranged, combined, and/or adapted in various ways without deviating from the scope of the present disclosure. The various steps are described, for example, as being performed by a computing system onboard an autonomous vehicle for example purposes. One or more portions could also, or alternatively, be performed by a system offboard/remote from the autonomous vehicle.

(152) The computer-implemented method **700** can include, at **702**, obtaining (e.g., by a computing system including one or more computing devices) scene data including one or more actor features. In some implementations, the scene data can be extracted or otherwise obtained from one or more scene observations. For instance, in some implementations, the scene data can be obtained by method **800** of FIG. 8.

(153) Additionally and/or alternatively, the method **700** can include, at **704**, providing (e.g., by the computing system) the scene data to a latent prior model. The latent prior model can be configured to generate scene latent data in response to receipt of scene data. In some implementations, the latent prior model can be or can include a scene interaction module including one or more graph neural networks. The scene latent data can be or can include one or more latent variables. In some implementations, the scene latent data can include a latent distribution that is partitioned into one or more latent variables. For instance, in some implementations, the one or more latent variables can be respective to the one or more actor features such that each actor feature has an associated latent variable of the scene latent data that is anchored to the actor feature. In some implementations, the one or more latent variables can be or can include one or more continuous latent variables. Additionally and/or alternatively, the method **700** can include, at **706**, obtaining (e.g., by the computing system) the scene latent data from the latent prior model.

(154) For instance, the generative process of future trajectories over actors can be formulated with a latent variable model including one or more latent variables that intuitively capture unobserved scene dynamics such as actor goals and style, multi-agent interactions, or future traffic light states. This modeling intuitively encourages the latent distribution to capture stochasticity in the generative process. In some implementations, the latent distribution can be a continuous latent distribution including one or more continuous latent variables for high expressivity.

(155) Producing a latent distribution that can capture all the uncertainties in any scenario can present challenges, as scenarios can vary drastically in the number of actors, the road topology, and/or traffic rules. This challenge can be mitigated by partitioning the scene latent distribution and obtaining a distributed representation where each latent variable in the scene latent distribution is anchored to a respective actor, such as anchored to a particular node in an interaction graph with traffic participants as nodes. The distributed representation may provide the benefit of naturally scaling the capacity of the latent distribution as the number of actors grows. Furthermore, the

anchoring may give the model an inductive bias that eases the learning of a scene latent distribution. Intuitively, each anchored latent variable encodes unobserved dynamics most relevant to its respective actor, including interactions with neighboring actors and traffic rules that apply in its locality. For instance, each latent variable can be represented as a diagonal multivariate Gaussian. Each latent variable can be conditioned on all actors such that the latent distribution is not marginally independent across actors, although factorized. For instance, despite anchoring each partition of the scene latent to an actor, each individual latent variable can contain information about the full scene, since each final node representation is dependent on the entire input because of the message propagation in the fully-connected interaction graph.

(156) Additionally and/or alternatively, the method **700** can include, at **708**, sampling (e.g., by the computing system) latent sample data from the scene latent data. For instance, the scene latent data (e.g., the one or more latent variables) can define a latent (e.g., a latent distribution) that can be sampled to produce latent samples of the scene latent data. The latent sample data can define a possible future for the actors (e.g., the actor features).

(157) Additionally and/or alternatively, the method **700** can include, at **710**, providing (e.g., by the computing system) the latent sample data to a decoder model. The decoder model can be configured to decode the latent sample data into a motion forecast including one or more predicted trajectories of the one or more actor features. For instance, the decoder model can produce a motion forecast from latent samples. Additionally and/or alternatively, the method **700** can include, at **712**, receiving (e.g., by the computing system) the motion forecast including one or more predicted trajectories of the one or more actor features from the decoder model.

(158) In some implementations, the decoder model can be or can include a deterministic decoder model. For instance, the decoder can be or can include a deterministic mapping to implicitly characterize the joint probability, such as opposed to explicitly representing it in a parametric form. This approach can provide for evaluation without factorizing the joint distribution and thus can sidestep the associated challenges. In particular, the deterministic decoder can be highly beneficial for generating socially consistent trajectories. In this framework, generating scene-consistent future trajectories across actors is simple and highly efficient, as it may be performed with only one stage of parallel sampling.

(159) In some implementations, the decoder model can be or can include a scene interaction module including one or more graph neural networks. For instance, the decoder including a scene interaction module can predict a realization of the future at the scene level via message parsing. As an example, each actor trajectory can be established with respect to samples from each latent variable of the scene latent data and/or each actor feature. This can provide for improved reasoning about multi-agent interactions such as car following, yielding, etc. For instance, each actor context can be initialized as a node in the decoder SIM. After a round of message parsing, each node can then contain an updated representation of a respective actor that takes into account the underlying dynamics of the scene summarized in the latent distribution. Finally, the trajectory sample for the actor can be deterministically decoded, such as by the output function of the SIM. For instance, the output function can be the deterministic mapping. This can be performed without requiring any additional sampling steps. The trajectory-level scene sample can thus be a collection of all actor trajectories.

(160) In some implementations, the decoder model can be or can include a specified and tractable conditional likelihood. In this implementation, many tools are available for inference and learning. As one example, variational inference, such as the variational auto-encoder (VAE), can be used.

(161) In some implementations, the decoder can be characterized via a stochastic sampling procedure where a conditional dependence is not specified. In this case, likelihood-free inference techniques can be required for learning. As examples, density estimation by comparison using either density ratio (GAN) or density difference (MMD). These techniques, however, may be more difficult to optimize.

(162) In some implementations the prior model, the encoder model, and/or the decoder model can include a scene interaction model. The scene interaction model can be configured to model the latent distribution as an interaction graph including one or more nodes representative of the one or more actor features and one or more edges representative of interactions between the one or more actor features. For instance, in some implementations, the scene interaction model can include one or more graph neural networks. In some implementations, a message function of the one or more graph neural networks can include a multi-layer perceptron model that takes as input one or more terminal nodes of the one or more nodes at a previous propagation step of the one or more graph neural networks. For instance, the edge or message function of the graph neural network(s) in the scene interaction module can include, for example, a 3-layer MLP (multi-layer perceptron) that takes as input the hidden states of the two terminal nodes at each edge in the graph at the previous propagation step. Additionally and/or alternatively, the input can include the projected coordinates of their corresponding bounding boxes. In some implementations, an aggregation function of the one or more graph neural networks includes a feature-wise max-pooling aggregation function. In some implementations, a gated recurrent unit cell is configured to update a state (e.g., hidden state) of the one or more nodes. In some implementations, the scene interaction module can include an output network that outputs the results from the graph propagations, such as a 2-layer MLP.

(163) The scene interaction module can model the latent distribution as an interaction graph, which can provide improved understanding of spatial information. This spatial information can be beneficial in jointly forecasting future trajectories of each actor. For instance, the node state of each node can be initialized with a set of actor features and known spatial information. The spatial information can include information such as relative coordinates of the actors relative to their peers or neighbors. In some cases, during object detection and local feature extraction around each actor, however, it may not be possible to include some necessary global information due to the limited receptive field and/or the translation invariance of convolutional neural networks. To remedy this, the node states can be initialized as the concatenation of the deep local features and the spatial information of each actor or node in the graph, such as its location, heading and/or its dimensions (e.g., in Bird's Eye View). A learned double edge function can propagate messages around the nodes in the graph. Given these messages, each actor can aggregate the messages (e.g., via max pooling) to update a respective node state. In some implementations, the scene interaction model can perform a single round of message passing to update the nodes' representation, taking into account spatiotemporal relationships. The scene interaction module in the prior, encoder and/or decoder can capture scene-level understanding that is not present with independence assumptions at the latent or output level.

(164) The systems and methods described herein can provide for efficiently sampling multiple possible futures from the latent distribution of the scene latent distribution. For instance, in some implementations, the method **700** can further include sampling (e.g., by the computing system) second latent sample data from the scene latent data. The second latent sample data can be descriptive of a differing possible future from the latent sample data. Additionally and/or alternatively, the method **700** can further include providing (e.g., by the computing system) the second latent sample data to the decoder model and receiving (e.g., by the computing system) a second motion forecast including one or more second predicted trajectories of the one or more actor features from the decoder model. The second predicted trajectories can differ from the predicted trajectories of other samples.

(165) For instance, a first sample from the scene latent distribution provided to the decoder can produce one possible realization of the future trajectories. A second sample can result in a distinct future. The sampling process is illustrated sequentially for the purposes of illustration, parallel sampling and/or decoding can be employed in accordance with example aspects of the present disclosure. For instance, the samples may be independent, as the stochasticity in the system is present in the latent distribution (e.g., as opposed to the decoder).

(166) FIG. 8 depicts a flow chart diagram of an example method **800** for recognizing actor features from sensor data according to example aspects of the present disclosure. One or more portions of the method **800** described herein can be implemented as an algorithm on the hardware components of the devices described herein (e.g., as in FIGS. 1, 2, 12, 13) to, for example, generate scenarios using parametric modeling and conducting simulations to test autonomous vehicle (e.g., autonomous vehicle **105** in FIG. 1) using those generated scenarios. Although FIG. 8 depicts steps performed in a particular order for purposes of illustration and discussion, method **800** of FIG. 8 is not limited to the particularly illustrated order or arrangement. The various steps of the methods disclosed herein can be omitted, rearranged, combined, and/or adapted in various ways without deviating from the scope of the present disclosure. The various steps are described, for example, as being performed by a computing system onboard an autonomous vehicle for example purposes. One or more portions could also, or alternatively, be performed by a system offboard/remote from the autonomous vehicle.

(167) For instance, the method **800** can include, at **802**, obtaining (e.g., by a computing system including one or more computing devices) one or more scene observations. The scene observations can be or can include data descriptive of sensor observations from one or more sensors configured to observe the scene, such as, for example, one or more sensors mounted on an autonomous vehicle. The sensors can be any suitable sensors, such as, for example, cameras, LIDAR sensors, etc. As one example, in some implementations, the scene observations can be and/or can include a three-dimensional (3D) LIDAR point cloud. In some implementations, the LIDAR data can be voxelized. In some implementations, the scene observations can be oriented in a “Birds-Eye View” (BEV) or top-down representation.

(168) In some implementations, the scene observations can additionally include map data, such as data descriptive of properties of roads, crosswalks, signage, intersections, railroads, buildings, and/or other terrain features of the scene. In some implementations, the map data can be rasterized. The map data can encode traffic elements such as intersections, lanes, roads, and traffic lights. In some implementations, elements with different semantics are encoded into different channels in the raster. Map elements that are rasterized can be or can include, for example, drivable surface polygons, road polygons, intersection polygons, straight vehicle lane polygons, dedicated left and right vehicle lane polygons, dedicated bike lane polygons, dedicated bus lane polygons, centerline markers for all lanes, lane dividers for all lanes with semantics (e.g., allowed to cross, not allowed to cross, might be allowed to cross).

(169) In some implementations, the height dimension of the sensor observations can be normalized with dense ground-height information provided by map data. In some implementations, multiple LiDAR sweeps can be used to exploit motion cues by compensating the ego-motion, such as by projecting the past sweeps to the coordinate frame of the current sweep). Furthermore, in some implementations, the height and time dimension is raveled into the channel dimension, to provide for the use of 2D convolution to process spatial and/or temporal information efficiently. The final representation may thus be a 3D occupancy tensor.

(170) The method **800** can include, at **804**, providing (e.g., by the computing system) the one or more scene observations to a scene feature extraction model. The scene feature extraction model can include one or more neural networks configured to produce scene data including one or more actor features from the one or more scene observations. For instance, in some cases, the scene features can be extracted from raw sensor data and HD maps in a differentiable manner, such that perception and motion forecasting can be trained jointly end-to-end. In some implementations, the scene feature extraction model can be or can include a backbone network. For instance, the backbone network can be a lightweight backbone network adapted for feature extraction. In some implementations, two separate streams can be instantiated such that the voxelized LiDAR and rasterized map are processed separately. The resulting features from both streams can then be concatenated feature-wise (e.g., if they share the same spatial resolution) and fused by a

convolutional header. These extracted features can inform both the downstream detection and motion forecasting networks. The method **800** can then include, at **806**, receiving (e.g., by the computing system) the one or more scene features from the scene feature extraction model. The scene features may be features that may not each correspond to an actor in the scene, such as a global feature map. For instance, a CNN-based perception backbone network architecture can be used to extract rich geometrical and motion features about the whole scene from a past history of voxelized LiDAR point clouds and/or a raster map.

(171) The method **800** can include, at **808**, providing (e.g., by the computing system) the one or more scene features to an actor feature recognition model. The actor feature recognition model can parse the scene features into actor features that correspond to an actor. For instance, the one or more actor features can include data descriptive of an actor context of one or more traffic participants. In some implementations, the actor feature recognition model can implement some or all of the method **900** of FIG. 9.

(172) For instance, in some implementations, the (e.g., scene) feature extraction model (e.g., the backbone model) includes two convolutional layers to output a classification or confidence score and/or a bounding box for each anchor location (e.g., each scene feature). These features are eventually reduced to the final set of candidates by applying non-maximal suppression (NMS) and finally thresholding low probability detections (given by the desired common recall). In some implementations, a backbone network along with features for object detection and per actor feature extraction are provided. The proposed mixture of trajectories output parameterization, where each way-point is a gaussian, is then used. In some cases, these baselines may not obtain temporally consistent samples, since the gaussians are independent across time (e.g., the models are not auto-regressive). To solve this, a heuristic sampler can be used to obtain temporally consistent samples from this model. The sampled trajectories are extracted using the re-parameterization technique for a bi-variate normal, where the model predicts a normal distribution per waypoint.

(173) In some cases, the noise can be constant across time for a given sample and actor. Intuitively, having a constant noise across time steps can provide sampled waypoints whose relative location with respect to its predicted mean and covariance is constant across time (e.g., translated by the predicted mean and scaled by the predicted covariance per time). In some cases, to address the compounding error problem found in some auto-regressive models, adjustments can be made to the training procedure to account for the noise in the conditioning space. To help simulate the noise it sees during inference, gaussian noise can be added to the conditioning state. The amount of noise expected between time-steps can be tuned.

(174) In some implementations, an object detection module can be responsible for recognizing other traffic participants in a scene, followed by a motion forecasting module that predicts how the scene might unroll given the current state or actor state of each actor. The actor state may be a compact representation of an actor, including qualities such as pose, velocity, and acceleration. This can be beneficial in some cases; however, it can be difficult to incorporate uncertainty due to sensor noise or occlusion. In some implementations, these two tasks can be combined by having a single model (e.g., a single fully convolutional backbone network) predict both the current and future states of actors. For instance, a single fully convolutional backbone network can predict both the current state and future state(s) for each pixel (e.g., in a bird's eye view grid) directly from sensor data, such as a voxelized LiDAR point-cloud, and/or map data, such as a semantic raster of an HD map. This approach can propagate uncertainty between the two tasks in the feature space, without the need of explicit intermediate representations.

(175) For instance, the perception and prediction tasks seek to understand where other actors are currently located and/or how they might move in the next few seconds. This can be accomplished by leveraging (e.g., 3D) sensor data such as LiDAR point clouds for dynamic object recognition and/or high-definition maps which provide information about the static part of the environment. For instance, scene features can be extracted from LiDAR and maps and subsequently fused to

produce rich features. Once rich features from the whole scene are extracted, object detection can be performed to recognize actor features corresponding to actors in the scenes.

(176) FIG. 9 depicts a flow chart diagram of an example method 900 for recognizing actor features from scene features according to example aspects of the present disclosure. One or more portions of the method 900 described herein can be implemented as an algorithm on the hardware components of the devices described herein (e.g., as in FIGS. 1, 2, 12, 13) to, for example, generate scenarios using parametric modeling and conducting simulations to test autonomous vehicle (e.g., autonomous vehicle 105 in FIG. 1) using those generated scenarios. Although FIG. 9 depicts steps performed in a particular order for purposes of illustration and discussion, method 900 of FIG. 9 is not limited to the particularly illustrated order or arrangement. The various steps of the methods disclosed herein can be omitted, rearranged, combined, and/or adapted in various ways without deviating from the scope of the present disclosure. The various steps are described, for example, as being performed by a computing system onboard an autonomous vehicle for example purposes. One or more portions could also, or alternatively, be performed by a system offboard/remote from the autonomous vehicle.

(177) For instance, the method 900 can include, at 902, obtaining (e.g., by a computing system including one or more computing devices) one or more scene observations. The scene observations can be or can include data descriptive of sensor observations from one or more sensors configured to observe the scene, such as, for example, one or more sensors mounted on an autonomous vehicle. The sensors can be any suitable sensors, such as, for example, cameras, LIDAR sensors, etc. As one example, in some implementations, the scene observations can be and/or can include a three-dimensional (3D) LIDAR point cloud. In some implementations, the LIDAR data can be voxelized. In some implementations, the scene observations can be oriented in a “Birds-Eye View” (BEV) or top-down representation.

(178) In some implementations, the scene observations can additionally include map data, such as data descriptive of properties of roads, crosswalks, signage, intersections, railroads, buildings, and/or other terrain features of the scene. In some implementations, the map data can be rasterized. The map data can encode traffic elements such as intersections, lanes, roads, and traffic lights. In some implementations, elements with different semantics are encoded into different channels in the raster. Map elements that are rasterized can be or can include, for example, drivable surface polygons, road polygons, intersection polygons, straight vehicle lane polygons, dedicated left and right vehicle lane polygons, dedicated bike lane polygons, dedicated bus lane polygons, centerline markers for all lanes, lane dividers for all lanes with semantics (e.g., allowed to cross, not allowed to cross, might be allowed to cross).

(179) In some implementations, the height dimension of the sensor observations can be normalized with dense ground-height information provided by map data. In some implementations, multiple LiDAR sweeps can be used to exploit motion cues by compensating the ego-motion, such as by projecting the past sweeps to the coordinate frame of the current sweep). Furthermore, in some implementations, the height and time dimension is raveled into the channel dimension, to provide for the use of 2D convolution to process spatial and/or temporal information efficiently. The final representation may thus be a 3D occupancy tensor.

(180) The method 900 can include, at 904, extracting (e.g., by a computing system) spatial feature maps for bounding boxes from the one or more scene features by rotated region of interest (ROI) align. Rotated ROI align can be applied to extract (e.g., fixed size) spatial feature maps for bounding boxes with arbitrary shapes and rotations from the scene features (e.g., the global feature map extracted by the backbone). For instance, rotated ROI align can provide actor contexts for each actor.

(181) The method 900 can include, at 906, pooling (e.g., by the computing system) a region around each spatial feature map to produce pooled actor features. For instance, a region around each actor in its frame can be pooled, such as with an axis defined by the actor's centroid orientation. The

pooled actor features may be larger than the eventual actor features.

(182) The method **900** can include, at **908**, downsampling (e.g., by the computing system) the pooled actor features by applying one or more downsampling convolutional neural networks. As one example, a 4-layer down-sampling convolutional network can be applied.

(183) The method **900** can include, at **910**, max-pooling (e.g., by the computing system) along spatial dimensions to reduce each pooled actor feature to a respective actor feature of the one or more actor features. For instance, after the downsampling CNN is applied, it can be followed by max-pooling along the spatial dimensions to reduce the feature map to a desired dimensional feature vector per actor. One example convolutional network uses a dilation factor of 2 for the convolutional layers to enlarge the receptive field for the per-actor features, which can improve performance.

(184) FIG. **10** depicts a flow chart diagram of an example method **1000** for training a motion forecasting system according to example aspects of the present disclosure. One or more portions of the method **1000** described herein can be implemented as an algorithm on the hardware components of the devices described herein (e.g., as in FIGS. **1**, **2**, **12**, **13**) to, for example, generate scenarios using parametric modeling and conducting simulations to test autonomous vehicle (e.g., autonomous vehicle **105** in FIG. **1**) using those generated scenarios. Although FIG. **10** depicts steps performed in a particular order for purposes of illustration and discussion, method **1000** of FIG. **10** is not limited to the particularly illustrated order or arrangement. The various steps of the methods disclosed herein can be omitted, rearranged, combined, and/or adapted in various ways without deviating from the scope of the present disclosure. The various steps are described, for example, as being performed by a computing system onboard an autonomous vehicle for example purposes. One or more portions could also, or alternatively, be performed by a system offboard/remote from the autonomous vehicle.

(185) The method **1000** can include, at **1002**, obtaining (e.g., by a computing system including one or more computing devices) a training dataset including one or more training examples labeled with ground truth data. For instance, the one or more training examples can include one or more actor features. Additionally and/or alternatively, the ground truth data can include a ground truth context of the one or more actor features. The ground truth context can be descriptive of a known context of the actor features, such as a known position, pose, velocity, etc.

(186) The method **1000** can include, at **1004**, providing (e.g., by the computing system) the one or more training examples labeled with ground truth data to a latent encoder model. The latent encoder model can be configured to produce a first latent distribution in response to receipt of the one or more training examples and the ground truth data. For instance, in some cases, integration over the scene latent distribution is intractable, in which case amortized variational inference can be used. For instance, by introducing an encoder distribution to approximate the true posterior, the learning problem can be reformulated as a maximization of the Evidence Lower BOund (ELBO). In some implementations, the latent encoder model can include a scene interaction module. For instance, after running one round of message passing, the scene interaction module can predict the distribution over latent variables.

(187) The encoder model can approximate the true posterior latent distribution. This model may also be called a recognition network. Because the encoder model receives the ground truth data (e.g., the target output) as an input, it can recognize scene dynamics that are unobserved by the latent prior model. In this manner, the encoder may only be used during training, since it requires access to the ground-truth future trajectories. For instance, the encoder may be omitted from deployed models and/or included at an online training system.

(188) The method **1000** can include, at **1006**, providing (e.g., by the computing system) the one or more training examples to a latent prior model. The latent prior model can be configured to produce a second latent distribution in response to receipt of the one or more training examples. For instance, the latent prior model can be agnostic to the ground truth data such that the latent

prior model is usable during inference (e.g., during prediction after ground truth data has already been used to train the latent prior model).

(189) The method **1000** can include, at **1008**, determining (e.g., by the computing system) a training loss based at least in part on the first latent distribution and the second latent distribution. The training loss can be or can include various losses from components of the model. For instance, in some implementations, the training loss can be or can include a KL divergence loss between the first latent distribution and the second latent distribution.

(190) Additionally and/or alternatively, in some implementations, the training loss can include a feature loss. The feature loss can be a loss in recognized features and/or actors. For instance, the feature loss can be determined by providing (e.g., by the computing system) training scene observations to a feature extraction model. The feature extraction model can be or can include a scene feature extraction model and/or an actor feature recognition model. Determining the feature loss can further include receiving (e.g., by the computing system) one or more predicted features from the feature extraction model. Determining the feature loss can further include determining (e.g., by the computing system) a feature loss between the one or more predicted features and the ground truth data. The feature loss can include a cross-entropy loss between the one or more predicted features and one or more training features of the ground truth data. Additionally and/or alternatively, the feature loss can include a regression loss between bounding boxes of the one or more predicted features and bounding boxes of the ground truth data. For instance, in some implementations, the feature loss includes a cross entropy loss applied to detection scores of the predicted features and/or a smooth L1 applied to the bounding box regression targets of the one or more predicted features.

(191) The method **1000** can include, at **1010**, backpropagating (e.g., by the computing system) the training loss through at least the latent prior model to train at least the latent prior model. For instance, the training loss can be backpropagated through some or all trainable portions of the implicit latent variable model, such as an object detection or feature extraction model, latent prior model, decoder model, encoder model, etc.

(192) In some implementations, the training loss can include a reconstruction loss. For instance, the reconstruction loss can be obtained by sampling (e.g., by the computing system) the first latent distribution to produce one or more first latent samples. The first latent distribution can include first sample data, such as a first set of samples from the latent distribution. Determining the reconstruction loss can further include sampling (e.g., by the computing system) the second latent distribution to produce one or more second latent samples. Determining the reconstruction loss can further include providing (e.g., by the computing system) the one or more first latent samples to a decoder model and receiving (e.g., by the computing system) one or more first predicted trajectories from the decoder model. The first predicted trajectories can be representative of a first possible future of the actor features. Determining the reconstruction loss can further include providing (e.g., by the computing system) the one or more second latent samples to the decoder model and receiving (e.g., by the computing system) one or more second predicted trajectories from the decoder model. The second predicted trajectories can be representative of a second possible future of the actor features that desirably matches the first possible future. Determining the reconstruction loss can include determining (e.g., by the computing system) a reconstruction loss based at least in part on the one or more first predicted trajectories and the one or more second predicted trajectories. For instance, the reconstruction loss can be based at least in part on a difference between the one or more first predicted trajectories and the one or more second predicted trajectories. In some implementations, the reconstruction loss can include a Huber loss.

(193) FIG. **11** depicts a flow chart diagram of an example method **1100** for controlling an autonomous vehicle according to example aspects of the present disclosure. One or more portions of the method **1100** described herein can be implemented as an algorithm on the hardware components of the devices described herein (e.g., as in FIGS. **1**, **2**, **12**, **13**) to, for example, generate

scenarios using parametric modeling and conducting simulations to test autonomous vehicle (e.g., autonomous vehicle **105** in FIG. **1**) using those generated scenarios. Although FIG. **11** depicts steps performed in a particular order for purposes of illustration and discussion, method **1100** of FIG. **11** is not limited to the particularly illustrated order or arrangement. The various steps of the methods disclosed herein can be omitted, rearranged, combined, and/or adapted in various ways without deviating from the scope of the present disclosure. The various steps are described, for example, as being performed by a computing system onboard an autonomous vehicle for example purposes. One or more portions could also, or alternatively, be performed by a system offboard/remote from the autonomous vehicle.

(194) The method **1100** can include, at **1102**, obtaining (e.g., by a computing system including one or more computing devices) one or more motion forecasts including one or more predicted trajectories of one or more actors features. For instance, the one or more predicted trajectories can be obtained from an implicit latent variable model according to example aspects of the present disclosure, such as by a computing device configured to implement any of methods **700**, **800**, **900**, or **1000** of FIGS. **7-10**.

(195) The method **1100** can include, at **1104**, providing (e.g., by the computing system) the one or more predicted trajectories to a motion planning model configured to generate a motion plan for an autonomous vehicle based at least in part on the one or more predicted trajectories. For instance, the motion planning model can generate a motion plan that safely navigates the scene in view of the predicted trajectories of the actors in the scene.

(196) The method **1100** can include, at **1106**, implementing (e.g., by the computing system) the motion plan to control the autonomous vehicle. For instance, one or more control devices (e.g., throttle, brakes, steering, etc.) can be controlled based at least in part on the motion plan to implement the motion plan and thus navigate the autonomous vehicle.

(197) FIG. **12** depicts a block diagram of a first example computing system **1200** according to example aspects of the present disclosure. Various means can be configured to perform the methods and processes described herein. FIG. **12** depicts example units associated with a computing system for performing operations and functions according to example embodiments of the present disclosure. As depicted, FIG. **12** depicts a computing system **1200** that can include, but is not limited to, scene observation obtaining unit(s) **1202**; feature extraction unit(s) **1204**; latent distribution modeling unit(s) **1206**; latent sampling unit(s) **1208**; decoding unit(s) **1210**; and motion planning unit(s) **1212**. In some implementations, one or more units may be implemented separately. In some implementations, one or more units may be included in one or more other units.

(198) The means can be programmed to perform one or more algorithm(s) for carrying out the operations and functions described herein (including the claims). For instance, the means can be configured to obtain one or more scene observations. For example, a computing system can obtain the scene observations from one or more sensors, one or more databases, etc. A scene observation obtaining unit **1202** is one example of a means for obtaining one or more scene observations.

(199) The means can be configured to extract one or more actor features from the one or more scene observations. For instance, the one or more scene observations can be provided to a feature extraction model. A feature extraction unit **1204** is one example of a means for extracting one or more actor features from the one or more scene observations.

(200) The means can be configured to model a latent distribution from the one or more actor features. For example, a latent prior model can be configured to produce a latent distribution, such as a distributed latent distribution, based on the one or more actor features. A latent distribution modeling unit **1206** is one example of a means for modeling the latent distribution.

(201) The means can be configured to sample the latent distribution to obtain latent sample data. For example, samples of the latent distribution can be obtained based at least in part on the latent distribution. A latent sampling unit **1208** is one example of a means for sampling the latent distribution.

(202) The means can be configured to decode the latent samples to produce a motion forecast including one or more predicted trajectories. For example, the latent sample data can be provided to a decoder model (e.g., a deterministic decoder) to produce the motion forecast. A decoding unit **1210** is one example of a means for decoding the latent samples.

(203) The means can be configured to determine a motion plan (e.g., for an autonomous vehicle. For example, a motion plan can be produced to allow an autonomous vehicle to safely and/or comfortably navigate a scene based at least in part on the motion forecast. A motion planning unit **1212** is one example of a means for determining a motion plan.

(204) In some implementations, one or more of the units may be implemented separately. In some implementations, one or more units may be a part of or included in one or more other units. These means can include processor(s), microprocessor(s), graphics processing unit(s), logic circuit(s), dedicated circuit(s), application-specific integrated circuit(s), programmable array logic, field-programmable gate array(s), controller(s), microcontroller(s), and/or other suitable hardware. The means can also, or alternately, include software control means implemented with a processor or logic circuitry, for example. The means can include or otherwise be able to access memory such as, for example, one or more non-transitory computer-readable storage media, such as random-access memory, read-only memory, electrically erasable programmable read-only memory, erasable programmable read-only memory, flash/other memory device(s), data registrar(s), database(s), and/or other suitable hardware.

(205) FIG. **13** depicts a block diagram of a second example computing system **1300** according to example aspects of the present disclosure. The example system **1300** includes a computing system **1302** and a machine learning computing system **1400** that are communicatively coupled over one or more networks **1380**. The computing system **1302** can include one or more computing devices **1305**.

(206) In some implementations, the computing device(s) **1305** can perform motion forecasting using an implicit latent variable model according to example aspects of the present disclosure. In some implementations, the computing device(s) **1305** can be included in an autonomous vehicle. For example, the computing device(s) **1305** can be on-board the autonomous vehicle. In other implementations, the computing device(s) **1305** is not located on-board the autonomous vehicle. For example, the computing device(s) **1305** can operate offline to predict future trajectories of one or more actors. The computing device(s) **1305** can include one or more distinct physical computing devices.

(207) The computing device(s) **1305** can include one or more processors **1310** and a memory **1315**. The one or more processors **1310** can be any suitable processing device (e.g., a processor core, a microprocessor, an ASIC, a FPGA, a controller, a microcontroller, etc.) and can be one processor or a plurality of processors that are operatively connected. The memory **1315** can include one or more non-transitory computer-readable storage media, such as RAM, ROM, EEPROM, EPROM, one or more memory devices, flash memory devices, etc., and combinations thereof.

(208) The memory **1315** can store information that can be accessed by the one or more processors **1310**. For instance, the memory **1315** (e.g., one or more non-transitory computer-readable storage mediums, memory devices) can store data **1320** that can be obtained, received, accessed, written, manipulated, created, and/or stored. The data **1320** can include, for instance, sensor data, map data, feature data, etc. as described herein. In some implementations, the computing system **1302** can obtain data from one or more memory device(s) that are remote from the computing system **1302**.

(209) The memory **1315** can also store computer-readable instructions **1325** that can be executed by the one or more processors **1320**. The instructions **1325** can be software written in any suitable programming language or can be implemented in hardware. Additionally, or alternatively, the instructions **1325** can be executed in logically and/or virtually separate threads on processor(s) **1310**.

(210) For example, the memory **1315** can store instructions **1325** that when executed by the one or

more processors **1310** cause the one or more processors **1310** (the computing system) to perform any of the operations and/or functions described herein, including, for example, obtaining scene data, obtaining scene latent data, etc.

(211) According to an aspect of the present disclosure, the computing device(s) **1305** can store or include one or more machine-learned models **1335**. As examples, the machine-learned models **1335** can be or can otherwise include various machine-learned models such as, for example, scene feature extraction models, actor feature recognition models, latent prior models, encoder models, and/or decoder models as described herein. Such machine-learned models **1335** can respectively include one or more neural networks (e.g., deep neural networks) including graph neural networks, support vector machines, decision trees, ensemble models, k-nearest neighbors models, Bayesian networks, or other types of models including linear models and/or non-linear models. Other example neural networks include feed-forward neural networks, recurrent neural networks (e.g., long short-term memory recurrent neural networks), convolutional neural networks, or other forms of neural networks.

(212) In some implementations, the computing system **1302** can receive the one or more machine-learned models **1335** from the machine learning computing system **1400** over network(s) **1380** and can store the one or more machine-learned models **1335** in the memory **1315**. The computing system **1302** can then use or otherwise implement the one or more machine-learned models **1335** (e.g., by processor(s) **1310**). In particular, the computing system **1302** can implement the machine learned model(s) **1335** to produce one or more predicted trajectories as described herein.

(213) The machine learning computing system **1400** can include one or more computing devices **1405**. The machine learning computing system **1400** can include one or more processors **1410** and a memory **1415**. The one or more processors **1410** can be any suitable processing device (e.g., a processor core, a microprocessor, an ASIC, a FPGA, a controller, a microcontroller, etc.) and can be one processor or a plurality of processors that are operatively connected. The memory **1415** can include one or more non-transitory computer-readable storage media, such as RAM, ROM, EEPROM, EPROM, one or more memory devices, flash memory devices, etc., and combinations thereof.

(214) The memory **1415** can store information that can be accessed by the one or more processors **1410**. For instance, the memory **1415** (e.g., one or more non-transitory computer-readable storage mediums, memory devices) can store data **1420** that can be obtained, received, accessed, written, manipulated, created, and/or stored. The data **1420** can include, for instance, training feature data including ground truth data as described herein. In some implementations, the machine learning computing system **1400** can obtain data from one or more memory device(s) that are remote from the machine learning computing system **1400**.

(215) The memory **1410** can also store computer-readable instructions **1425** that can be executed by the one or more processors **1410**. The instructions **1425** can be software written in any suitable programming language or can be implemented in hardware. Additionally, or alternatively, the instructions **1425** can be executed in logically and/or virtually separate threads on processor(s) **1410**.

(216) For example, the memory **1415** can store instructions **1425** that when executed by the one or more processors **1410** cause the one or more processors **1410** (the computing system) to perform any of the operations and/or functions described herein, including, for example, determining a feature loss and/or backpropagating the feature loss.

(217) In some implementations, the machine learning computing system **1400** includes one or more server computing devices. If the machine learning computing system **1400** includes multiple server computing devices, such server computing devices can operate according to various computing architectures, including, for example, sequential computing architectures, parallel computing architectures, or some combination thereof.

(218) In addition or alternatively to the model(s) **1335** at the computing system **1302**, the machine

learning computing system **1400** can include one or more machine-learned models **1435**. As examples, the machine-learned models **1435** can be or can otherwise include various machine-learned models such as, for example, scene feature extraction models, actor feature recognition models, latent prior models, encoder models, and/or decoder models as described herein. Such machine-learned models **1435** can respectively include one or more neural networks (e.g., deep neural networks) including graph neural networks, support vector machines, decision trees, ensemble models, k-nearest neighbors models, Bayesian networks, or other types of models including linear models and/or non-linear models. Other example neural networks include feed-forward neural networks, recurrent neural networks (e.g., long short-term memory recurrent neural networks), convolutional neural networks, or other forms of neural networks.

(219) As an example, the machine learning computing system **1400** can communicate with the computing system **1302** according to a client-server relationship. For example, the machine learning computing system **1400** can implement the machine-learned models **1435** to provide a web service to the computing system **1302**.

(220) Thus, machine-learned models **1335** can be located and used at the computing system **1302** and/or machine-learned models **1435** can be located and used at the machine learning computing system **1400**.

(221) In some implementations, the machine learning computing system **1400** and/or the computing system **1302** can train the machine-learned models **1335** and/or **1340** through use of a model trainer **1440**. The model trainer **1440** can train the machine-learned models **1335** and/or **1440** using one or more training or learning algorithms. One example training technique is backwards propagation of errors. In some implementations, the model trainer **1440** can perform supervised training techniques using a set of labeled training data. In other implementations, the model trainer **1440** can perform unsupervised training techniques using a set of unlabeled training data. The model trainer **1440** can perform a number of generalization techniques to improve the generalization capability of the models being trained. Generalization techniques include weight decays, dropouts, or other techniques.

(222) In particular, the model trainer **1440** can train a machine-learned model **1335** and/or **1435** based on a set of training data **1445**. The training data **1445** can include, for example, training features (e.g., actor features) having known ground truth information stored in ground truth data. The model trainer **1440** can be implemented in hardware, firmware, and/or software controlling one or more processors.

(223) The computing system **1302** and the machine learning computing system **1400** can each include a communication interface **1330** and **1450**, respectively. The communication interfaces **1330/1450** can be used to communicate with one or more systems or devices, including systems or devices that are remotely located from the computing system **1302** and the machine learning computing system **1400**. A communication interface **1330/1450** can include any circuits, components, software, etc. for communicating with one or more networks (e.g., **1380**). In some implementations, a communication interface **1330/1450** can include, for example, one or more of a communications controller, receiver, transceiver, transmitter, port, conductors, software and/or hardware for communicating data.

(224) The network(s) **1380** can be any type of network or combination of networks that allows for communication between devices. In some embodiments, the network(s) can include one or more of a local area network, wide area network, the Internet, secure network, cellular network, mesh network, peer-to-peer communication link and/or some combination thereof and can include any number of wired or wireless links. Communication over the network(s) **1380** can be accomplished, for instance, via a network interface using any type of protocol, protection scheme, encoding, format, packaging, etc.

(225) FIG. **13** illustrates one example computing system **1300** that can be used to implement the present disclosure. Other computing systems can be used as well. For example, in some

implementations, the computing system **1302** can include the model trainer **1440** and the training dataset **1445**. In such implementations, the machine-learned models **1440** can be both trained and used locally at the computing system **1302**. As another example, in some implementations, the computing system **1302** is not connected to other computing systems.

(226) In addition, components illustrated and/or discussed as being included in one of the computing systems **1302** or **1400** can instead be included in another of the computing systems **1302** or **1400**. Such configurations can be implemented without deviating from the scope of the present disclosure. The use of computer-based systems allows for a great variety of possible configurations, combinations, and divisions of tasks and functionality between and among components. Computer-implemented operations can be performed on a single component or across multiple components. Computer-implemented tasks and/or operations can be performed sequentially or in parallel. Data and instructions can be stored in a single memory device or across multiple memory devices.

(227) Computing tasks discussed herein as being performed at computing device(s) remote from the autonomous vehicle can instead be performed at the autonomous vehicle (e.g., via the vehicle computing system), or vice versa. Such configurations can be implemented without deviating from the scope of the present disclosure. The use of computer-based systems allows for a great variety of possible configurations, combinations, and divisions of tasks and functionality between and among components. Computer-implemented operations can be performed on a single component or across multiple components. Computer-implements tasks and/or operations can be performed sequentially or in parallel. Data and instructions can be stored in a single memory device or across multiple memory devices.

(228) Aspects of the disclosure have been described in terms of illustrative embodiments thereof. Numerous other embodiments, modifications, and/or variations within the scope and spirit of the appended claims can occur to persons of ordinary skill in the art from a review of this disclosure. Any and all features in the following claims can be combined and/or rearranged in any way possible.

(229) While the present subject matter has been described in detail with respect to various specific example embodiments thereof, each example is provided by way of explanation, not limitation of the disclosure. Those skilled in the art, upon attaining an understanding of the foregoing, can readily produce alterations to, variations of, and/or equivalents to such embodiments. Accordingly, the subject disclosure does not preclude inclusion of such modifications, variations, and/or additions to the present subject matter as would be readily apparent to one of ordinary skill in the art. For instance, features illustrated and/or described as part of one embodiment can be used with another embodiment to yield a still further embodiment. Thus, it is intended that the present disclosure cover such alterations, variations, and/or equivalents.

Claims

1. A computer-implemented method for determining scene-consistent motion forecasts from sensor data, the method comprising: obtaining, by a computing system, scene latent data comprising a latent distribution from a scene interaction module, the latent distribution encoding dynamics relative to a respective actor of one or more actors; wherein the scene interaction module is configured to model the latent distribution as an interaction graph, wherein the interaction graph comprises one or more nodes representative of the one or more actors, and wherein the latent distribution is partitioned into one or more latent variables; updating, by the computing system, one or more node states respectively associated with the one or more nodes of the interaction graph based on spatiotemporal information of the scene latent data respective of the one or more actors; obtaining, by the computing system, second scene latent data from the scene interaction module comprising the updated one or more node states respectively associated with the one or more nodes of the interaction graph; receiving, by the computing system, an output associated with a motion

- forecast comprising one or more predicted trajectories of the one or more actors based on the second scene latent data; providing, by the computing system, the one or more predicted trajectories to a motion planning model configured to generate a motion plan for an autonomous vehicle based at least in part on the one or more predicted trajectories; and implementing, by the computing system, the motion plan to control the autonomous vehicle.
2. The computer-implemented method of claim 1, wherein the latent distribution is partitioned into one or more latent variables, wherein the one or more latent variables of the latent distribution are anchored to respective actors of the one or more actors.
 3. The computer-implemented method of claim 1, wherein the scene interaction module comprises one or more graph neural networks.
 4. The computer-implemented method of claim 3, wherein the scene latent data encodes interactions between the one or more actors by the one or more graph neural networks.
 5. The computer-implemented method of claim 1, wherein the interaction graph comprises one or more edges between the one or more nodes, the one or more edges representative of interactions between the one or more actors respective to the one or more nodes.
 6. The computer-implemented method of claim 5, wherein updating, by the computing system, the one or more node states respectively associated with the one or more nodes of the interaction graph comprises passing one or more messages along the one or more edges.
 7. The computer-implemented method of claim 6, wherein updating, by the computing system, the one or more node states comprises aggregating the messages between the one or more nodes to update respective node states of the one or more node states.
 8. The computer-implemented method of claim 1, further comprising updating, by the computing system, the one or more node states to a final node state, wherein the final node state is indicative of a full scene.
 9. The computer-implemented method of claim 1, further comprising initializing, by the computing system, respective node states with spatial information, the spatial information indicative of a location of an actor relative to other actors.
 10. The computer-implemented method of claim 1, wherein obtaining, by the computing system, the scene latent data comprises: obtaining scene data comprising one or more actor features; providing the scene data to the scene interaction module; and receiving the scene latent data as an output of the scene interaction module.
 11. A computing system comprising: one or more processors; and one or more non-transitory computer-readable media storing instructions executable by the one or more processors to perform operations, the operations comprising: obtaining scene latent data comprising a latent distribution from a scene interaction module, the latent distribution encoding dynamics relative to a respective actor of one or more actors; wherein the scene interaction module is configured to model the latent distribution as an interaction graph, wherein the interaction graph comprises one or more nodes representative of the one or more actors, and wherein the latent distribution is partitioned into one or more latent variables; updating one or more node states respectively associated with the one or more nodes of the interaction graph based on spatiotemporal information of the scene latent data respective of the one or more actors; obtaining second scene latent data from the scene interaction module comprising the updated one or more node states respectively associated with the one or more nodes of the interaction graph; receiving an output associated with a motion forecast comprising one or more predicted trajectories of the one or more actors based on the second scene latent data; providing the one or more predicted trajectories to a motion planning model configured to generate a motion plan for an autonomous vehicle based at least in part on the one or more predicted trajectories; and implementing the motion plan to control the autonomous vehicle.
 12. The computing system of claim 11, wherein the latent distribution is partitioned into one or more latent variables, wherein the one or more latent variables of the latent distribution are anchored to respective actors of the one or more actors.

13. The computing system of claim 11, wherein the scene interaction module comprises one or more graph neural networks, and wherein the scene latent data encodes interactions between the one or more actors by the one or more graph neural networks.
 14. The computing system of claim 11, wherein the interaction graph comprises one or more edges between the one or more nodes, the one or more edges representative of interactions between the one or more actors respective to the one or more nodes.
 15. The computing system of claim 14, wherein updating the one or more node states respectively associated with the one or more nodes of the interaction graph comprises passing one or more messages along the one or more edges.
 16. The computing system of claim 15, wherein updating the one or more node states comprises aggregating the messages between the one or more nodes to update respective node states of the one or more node states.
 17. The computing system of claim 11, wherein the operations further comprise updating the one or more node states to a final node state, wherein the final node state is indicative of a full scene.
 18. The computing system of claim 11, wherein the operations further comprise initializing respective node states with spatial information, the spatial information indicative of a location of an actor relative to other actors.
 19. The computing system of claim 11, wherein obtaining the scene latent data comprises: obtaining scene data comprising one or more actor features; providing the scene data to the scene interaction module; and receiving the scene latent data as an output of the scene interaction module.
 20. A non-transitory computer-readable media storing instructions executable by one or more processor to cause the processors to perform operations, the operations comprising: obtaining scene latent data comprising a latent distribution from a scene interaction module, the latent distribution encoding dynamics relative to a respective actor of one or more actors; wherein the scene interaction module is configured to model the latent distribution as an interaction graph, wherein the interaction graph comprises one or more nodes representative of the one or more actors, and wherein the latent distribution is partitioned into one or more latent variables; updating one or more node states respectively associated with the one or more nodes of the interaction graph based on spatiotemporal information of the scene latent data respective of the one or more actors; obtaining second scene latent data from the scene interaction module comprising the updated one or more node states respectively associated with the one or more nodes of the interaction graph; receiving an output associated with a motion forecast comprising one or more predicted trajectories of the one or more actors based on the second scene latent data; providing the one or more predicted trajectories to a motion planning model configured to generate a motion plan for an autonomous vehicle based at least in part on the one or more predicted trajectories; and implementing the motion plan to control the autonomous vehicle.
-