



US012387715B2

(12) **United States Patent**
Aggarwal et al.

(10) **Patent No.:** **US 12,387,715 B2**
(45) **Date of Patent:** ***Aug. 12, 2025**

(54) **LEARNING OFFLINE VOICE COMMANDS
BASED ON USAGE OF ONLINE VOICE
COMMANDS**

(58) **Field of Classification Search**
CPC G10L 15/22; G10L 15/30; G10L 15/063;
G10L 2015/223; G10L 15/26;
(Continued)

(71) Applicant: **GOOGLE LLC**, Mountain View, CA
(US)

(56) **References Cited**

(72) Inventors: **Vikram Aggarwal**, Palo Alto, CA (US);
Moises Morgenstern Gali, San
Francisco, CA (US)

U.S. PATENT DOCUMENTS

6,195,641 B1 2/2001 Loring et al.
6,873,951 B1 3/2005 Lin et al.

(Continued)

(73) Assignee: **GOOGLE LLC**, Mountain View, CA
(US)

FOREIGN PATENT DOCUMENTS

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 0 days.

CN 101075434 11/2007
CN 101454775 6/2009

(Continued)

This patent is subject to a terminal dis-
claimer.

OTHER PUBLICATIONS

(21) Appl. No.: **18/380,538**

China National Intellectual Property Administration; Notification of
Grant issued in Application No. 201980011141.4; 6pages; dated
Nov. 26, 2023.

(22) Filed: **Oct. 16, 2023**

(Continued)

(65) **Prior Publication Data**

US 2024/0038219 A1 Feb. 1, 2024

Primary Examiner — Olujimi A Adesanya

(74) *Attorney, Agent, or Firm* — Gray Ice Higdon

Related U.S. Application Data

(63) Continuation of application No. 17/521,472, filed on
Nov. 8, 2021, now Pat. No. 11,790,890, which is a
(Continued)

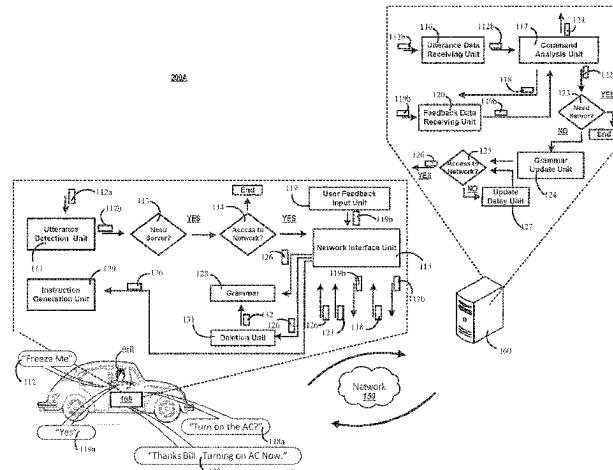
(51) **Int. Cl.**
G10L 21/00 (2013.01)
G10L 15/07 (2013.01)
(Continued)

(52) **U.S. Cl.**
CPC **G10L 15/07** (2013.01); **G10L 15/19**
(2013.01); **G10L 15/22** (2013.01); **G10L 15/30**
(2013.01);
(Continued)

(57) **ABSTRACT**

Methods, systems, apparatus, including computer programs
encoded on a computer storage medium, for a user device to
learn offline voice actions. In one aspect, the method
includes actions of detecting, by the user device, an utter-
ance at a first time when the user device is connected to a
server by a network, providing, by the user device, the
utterance to the server using the network, receiving, by the
user device and from the server, an update to the grammar
of the user device, detecting, by the user device, a subse-
quent utterance of the utterance at a second time when the
user device is not connected to the server by a network, and
in response to detecting, by the user device, the subsequent
utterance of the utterance at the second time, identifying, by

(Continued)



the user device, an operation to perform based on (i) the subsequent utterance, and (ii) the updated grammar.

20 Claims, 5 Drawing Sheets

Related U.S. Application Data

continuation of application No. 15/862,615, filed on Jan. 4, 2018, now Pat. No. 11,170,762.

(51) Int. Cl.

G10L 15/19 (2013.01)

G10L 15/22 (2006.01)

G10L 15/30 (2013.01)

G10L 15/32 (2013.01)

G10L 25/00 (2013.01)

(52) U.S. Cl.

CPC **G10L 15/32** (2013.01); **G10L 2015/228** (2013.01)

(58) Field of Classification Search

CPC G10L 15/065; G10L 15/183; G10L 2015/0635; G10L 15/00; G10L 17/00; G10L 15/20; G10L 15/18; G10L 15/07; G10L 15/28; G10L 17/04; G10L 2015/0631

See application file for complete search history.

(56)

References Cited

U.S. PATENT DOCUMENTS

8,185,392 B1	5/2012	Strope et al.	
8,370,146 B1	2/2013	Schalkwyk et al.	
8,805,684 B1 *	8/2014	Aleksic	G10L 15/07 704/270.1
9,070,367 B1	6/2015	Hoffmeister et al.	
9,836,527 B2	12/2017	Gao et al.	
10,109,273 B1 *	10/2018	Rajasekaram	G10L 15/183
11,170,762 B2	11/2021	Aggarwal et al.	
2002/0077823 A1	6/2002	Fox et al.	
2003/0046074 A1	3/2003	Ruback et al.	
2005/0096753 A1	5/2005	Arling et al.	
2005/0105712 A1	5/2005	Williams et al.	
2007/0005360 A1	1/2007	Huning et al.	
2007/0073540 A1	3/2007	Hirakawa et al.	
2007/0276651 A1	11/2007	Bliss et al.	
2008/0235021 A1	9/2008	Cross et al.	
2009/0112605 A1	4/2009	Gupta	
2010/0179812 A1	7/2010	Jang et al.	
2012/0179471 A1	7/2012	Newman et al.	
2013/0073293 A1 *	3/2013	Jang	G10L 15/22 704/E21.001
2013/0088419 A1	4/2013	Kim et al.	
2013/0132089 A1	5/2013	Fanty et al.	
2013/0204967 A1 *	8/2013	Seo	H04L 67/63 709/224

2013/0211824 A1 *	8/2013	Tseng	G06F 40/274 704/10
2014/0136013 A1	5/2014	Wolverton et al.	
2014/0200737 A1	7/2014	Lortz et al.	
2014/0244259 A1	8/2014	Rosario et al.	
2014/0274203 A1	9/2014	Ganong, III et al.	
2014/0337007 A1	11/2014	Waibel et al.	
2015/0045988 A1	2/2015	Gusikhin et al.	
2015/0255068 A1 *	9/2015	Kim	G10L 17/04 704/246
2015/0281401 A1	10/2015	Le et al.	
2015/0371628 A1 *	12/2015	Kreifeldt	G10L 15/02 704/254
2016/0094509 A1	3/2016	Ye	
2016/0109322 A1 *	4/2016	Eguchi	G01N 29/036 73/40
2016/0155443 A1	6/2016	Khan et al.	
2016/0179462 A1	6/2016	Bjorkengren	
2016/0252972 A1 *	9/2016	Kim	G10L 15/26 704/9
2016/0269524 A1	9/2016	Stottlemeyer	
2016/0322044 A1 *	11/2016	Jung	G10L 15/065
2016/0350320 A1	12/2016	Sung et al.	
2017/0069312 A1	3/2017	Sundararajan et al.	
2017/0162204 A1	6/2017	Wilson et al.	
2017/0242914 A1 *	8/2017	Gao	G10L 15/26
2018/0122365 A1	5/2018	Kunitake et al.	
2018/0174580 A1	6/2018	Kim et al.	
2018/0217810 A1	8/2018	Agrawal	
2019/0019516 A1	1/2019	Lawrence et al.	
2019/0034405 A1 *	1/2019	Ivan	G06F 40/279
2019/0108832 A1	4/2019	Tomar et al.	
2019/0206388 A1	7/2019	Aggarwal et al.	
2020/0394465 A1 *	12/2020	Gil Ramos	G06F 18/2155
2022/0157318 A1 *	5/2022	Sharifi	G10L 15/22

FOREIGN PATENT DOCUMENTS

CN	101588417	11/2009
CN	107093425	8/2017
WO	2016191319	12/2016

OTHER PUBLICATIONS

European Patent Office, Intention to Grant issued in Application No. 19701422.8, 44 pages, dated May 28, 2024.

European Patent Office; Communication pursuant Article 94(3) EPC issued in Application No. 19701422.8; 4 pages; dated Jul. 6, 2023.

China National Intellectual Property Administration; Notification of Second Office Action issued in Application No. 201980011141.4; 7 pages; dated Jul. 20, 2023.

China National Intellectual Property Administration; Notification of First Office Action issued in Application No. 201980011141.4; 30 pages; dated Dec. 5, 2022.

European Patent Office; Communication issued in Application No. 19701422.8, 4 pages, dated Oct. 20, 2022.

European Patent Office; International Search Report issued in Application No. PCT/US2019/012399; 14 pages; dated mailed Apr. 11, 2019.

* cited by examiner





FIG. 1



200B




FIG. 2B

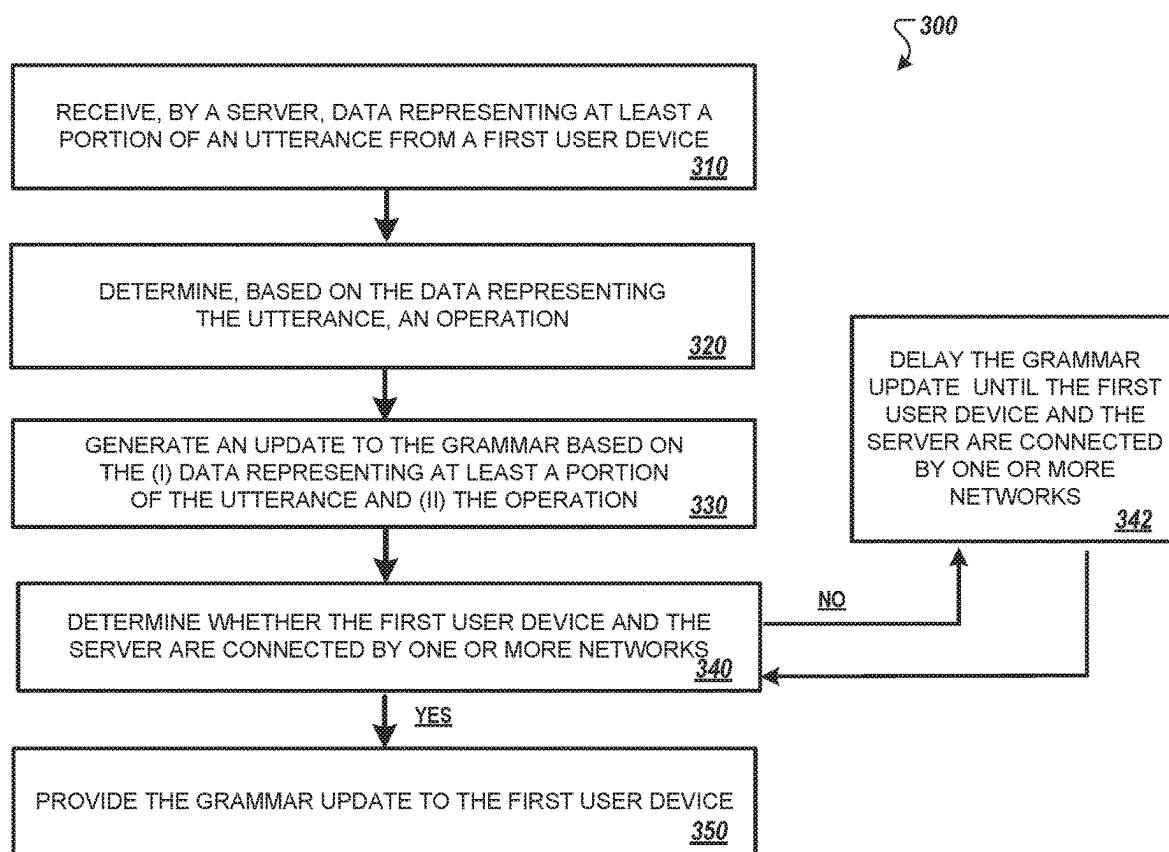


Fig. 3

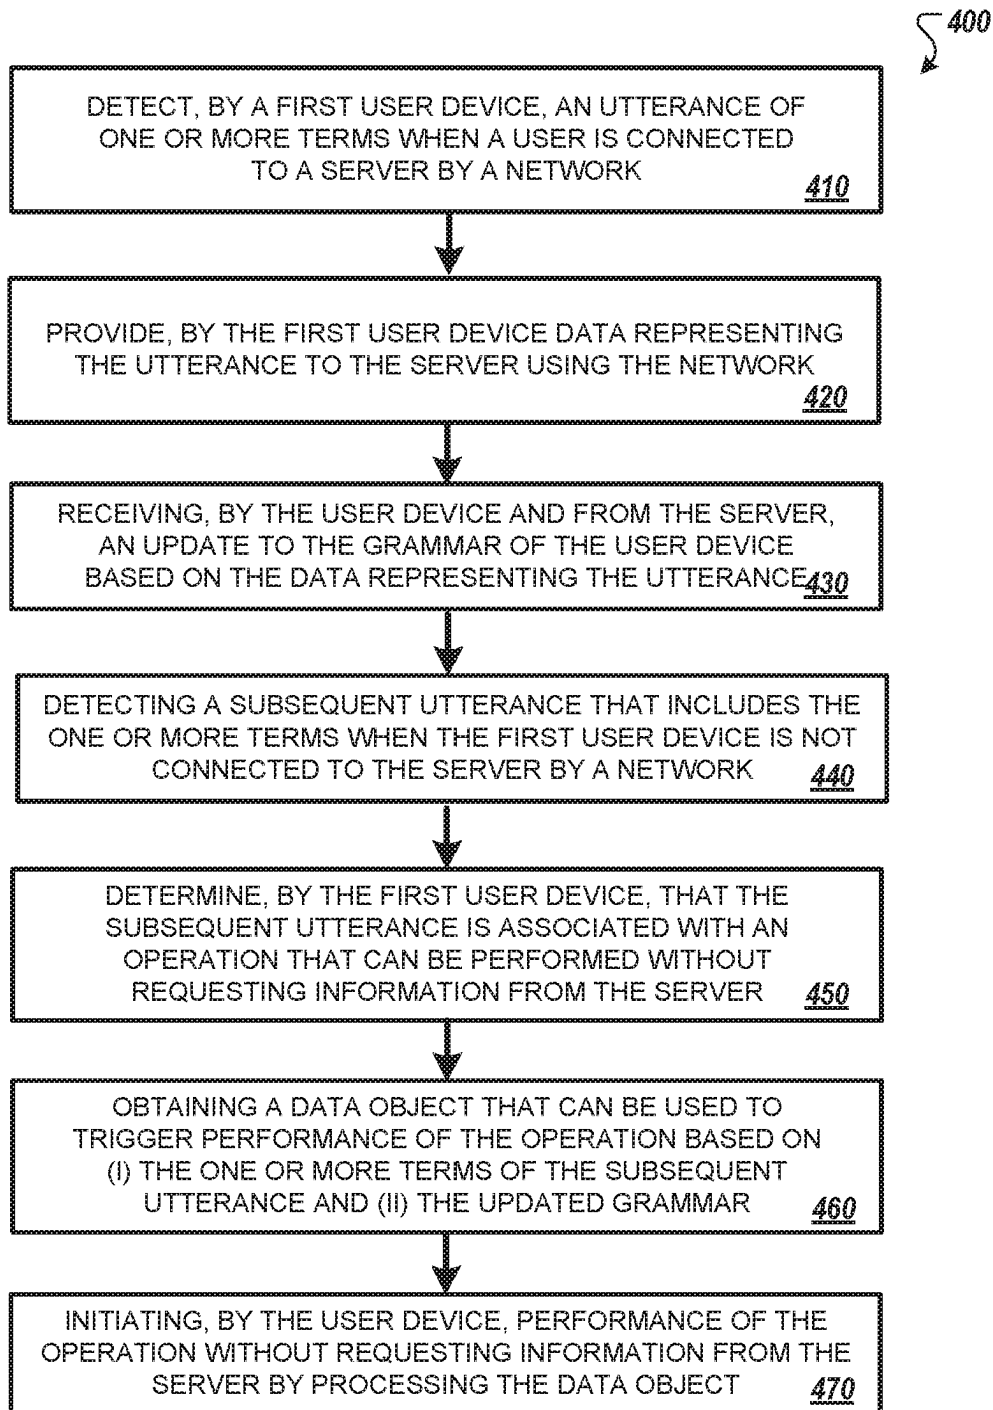


Fig. 4

1

LEARNING OFFLINE VOICE COMMANDS BASED ON USAGE OF ONLINE VOICE COMMANDS

BACKGROUND

People are relying on their user devices to perform a variety of different tasks. User devices may include, for example, a smartphone, a smartwatch, a tablet computer, a smart speaker, a headset, a car, a television, a thermostat, or the like. Such devices may include hardware and software that are configured to enable digital assistant functionality. Such applications are generally referred to as “digital assistants.”

Digital assistant functionality may include using a speech recognizer to detect an utterance that is uttered by a user and using the grammar of the speech recognizer to identify an operation that should be performed based on the detected utterance. Such assistant functionality can be generally broken down into online and offline functionality. Online functionality may include commands such as “what is the weather like in Paris today?,” “Order a large pizza for me,” or the like. Such commands require an active network connection and information from one or more servers access via the network.

In many systems, offline functionality is limited to a subset of commands that do not require network connectivity. For example, offline digital assistant functionality may be limited to a subset of commands such as “play 101.5 FM,” “increase volume,” “turn on air conditioner,” or the like.

SUMMARY

The present disclosure is directed towards a system and method for updating the grammar of a speech recognizer that is stored on a user device and used to execute offline voice commands. In general, one innovative aspect of the subject matter described in this specification can be embodied in methods that include the actions of detecting, by a user device, an utterance at a first time when the user device is connected to a server by a network, wherein the utterance includes one or more terms, providing, by the user device, the utterance to the server using the network, receiving, by the user device and from the server, an update to the grammar of the user device based on the one or more terms, detecting, by the user device, a subsequent utterance at a second time when the user device is not connected to the server by a network, wherein the subsequent utterance includes the one or more terms, and in response to detecting, by the user device, the subsequent utterance at the second time: determining, by the user device, that the subsequent utterance is associated with an operation that can be performed without requesting information from the server, identifying, by the user device, an operation to perform based on (i) the one or more terms of the subsequent utterance, and (ii) the updated grammar, and initiating, by the user device, performance of the operation without requesting information from the server. Other embodiments of this aspect include corresponding systems, apparatus, and computer programs, configured to perform the actions of the methods, encoded on computer storage devices.

These and other versions may optionally include one or more of the following features. For example, in some implementations, initiating by the user device, performance of the operation without requesting information from the server may include generating, by the user device, an action

2

comprising a data object that, when executed, causes the user device to initiate performance of an operation responsive to the subsequent utterance, and executing, by the user device, the action to trigger performance of the operation without requesting information from the server.

In some implementations, the method may also include actions of receiving, by the user device, an action. The action may include a data object that, when executed, causes the user device to initiate performance of an operation responsive to the subsequent utterance, wherein the action was generated by the server, and storing, by the user device, the action on the user device.

In some implementations, initiating by the user device, performance of the operation without requesting information from the server may include accessing, by the user device, the action that is stored on the user device, and executing, by the user device, the action to trigger performance of the operation without requesting information from the server.

In some implementations, receiving, by the user device and from the server, an update to the grammar of the user device based on the one or more terms may include receiving, by the user device, the one or more terms, and updating the grammar of the user device to include the one or more terms.

In some implementations, the one or more terms are associated with a threshold amount of time that indicates a maximum period of time the one or more terms will be included in the grammar absent the user device detecting the subsequent utterance of the one or more terms by a user.

In some implementations, the method may also include actions of determining, by the user device, that the threshold period of time has been exceeded without the user device detecting the subsequent utterance of the one or more terms by the user, and removing, by the user device, the one or more terms from the grammar of the user device.

Another innovative aspect of the subject matter described in this specification can be embodied in methods that include the actions of processing, by a server, a transcription of an utterance that was detected by a first user device when the server and the first user device are connected by a network, determining, by the server based on the processing of the transcription of the utterance, an operation that can be performed by the first user device without requesting information from the server, wherein the transcription includes one or more terms, and providing, by the server, an update to the grammar of the first user device based on the one or more terms. Other embodiments of this aspect include corresponding systems, apparatus, and computer programs, configured to perform the actions of the methods, encoded on computer storage devices.

These and other versions may optionally include one or more of the following features. For example, in some implementations, providing, by the server, an update to the grammar of the first user device based on the one or more terms may include providing the one or more terms to the first user device.

In some implementations, the method may also include actions of associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent the subsequent utterance of the one or more terms by a user to perform the operation.

In some implementations, determining, by the server based on a transcription of the utterance, an operation that can be performed by the first user device without requesting

3

information from the server may include generating, by the server, a data object that, when executed, causes the first user device to initiate performance of an operation responsive to the received utterance.

In some implementations, providing, by the server, the update to the grammar of the first user device, based on the one or more terms may include providing the generated data object to the first user device.

In some implementations, the update to the grammar enables the first user device to execute a data object that initiates performance of the operation in response to a subsequent receipt of the utterance without requesting information from the server.

In some implementations, the method may also include actions of identifying, by the server, a user profile that is associated with the first user device, identifying, by the server based on the user profile, a second user device that is registered to a user of the first user device, and providing, by the server, an update to the grammar of the second user device based on the one or more terms.

Particular embodiments of the subject matter described in this specification can be implemented so as to realize one or more of the following advantages. Updating the grammar of a speech recognizer using the system and method of the present disclosure enables the speech recognizer to learn personalized commands that can be used to execute offline voice commands. This provides the advantage of enabling a user to command a digital assistant using the user's own personalized language. If a user can utter commands that can be detected by a speech recognizer in the user's own personalized language and get positive results from applications that receive input commands detected by the speech recognizer, this will increase the likelihood that a user will use applications that can encourage voice input such as digital assistant applications relative to the use of such applications when a user is required memorize a set of standard commands made available to all users of the digital assistant.

The present disclosure provides another advantage in that the learning of personalized commands is achieved as a result of automated processes employed the systems and methods of the present disclosure. That is, a user does not need to initiate the personalization processes. Instead, the personalization of the speech recognizer's grammar can be achieved based on the normal operation of the system based in part on user commands and system responses.

In addition to encouraging a user to use applications that receive input commands detected by the speech recognizer, the present disclosure also improves the functionality of a computer. The updating of the grammar of the speech recognizer that is stored on the user device allows for an increase in the number commands that can be detected by the speech recognizer and executing using one or more control systems employed by the user device. As such, this require less consultation with a server accessible via network to disambiguate a previously unknown voice command, prompt a user for feedback that is used to help disambiguate the unknown voice command, and then receive data object from the server that can be processed by one or more control systems on the user device to execute a particular action (e.g., turn on air conditioner, turn off air conditioner, tune radio to a particular radio station, increase volume, turn off lights, or the like). Less consultation with a networked server results in the usage of less network bandwidth and a reduction in latency that is involved to execute a command from the time the command is initially uttered.

4

Yet other advantages may be achieved by the present disclosure. For example, a user device such as a phone, entertainment and climate system of an automobile, or the like, will have to contact a server less to disambiguate uttered commands once the local grammar has been updated. As a result, the user device has to communicate with one or more networks less, resulting in battery savings. By way of another example, a user device utilizing the present disclosure has to use less CPU resources and less storage space storing additional offline Language Packs. In some implementations, this may reduce the amount of downloads and storage of language packs to as little as 40 Mb downloads of Language pack. By way of another example, the server itself has to do less processing to create these large Language packs, which can have a cumulative effect as more languages are being added to grammar of a speech recognizer.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 is a contextual diagram of an example of a system for executing an offline voice command.

FIG. 2A is a contextual diagram of an example of a system for learning offline voice commands

FIG. 2B is a contextual diagram of the system of FIG. 2A that has learned an offline voice command.

FIG. 3 is a flowchart of an example of a process for generating a personalized grammar update for a user device.

FIG. 4 is a flowchart of an example of a process for personalizing a grammar of a user device.

DETAILED DESCRIPTION

FIG. 1 is a contextual diagram of an example of a system **100** for executing an offline voice command **112**. The system **100** includes a user device **105**, a speech recognizer **110**, an instruction generation unit **129**, and one or more control systems **131**.

The user device **105** can include any device that includes one or more processors to and one or more memory units storing instructions that, when executed by the one or more processors, realize the functionality of a speech recognizer **110**, an instruction generation unit **129**, and one or more control system **131**, as described herein. In the example of FIG. 1, the user device **105** is a computing system of a vehicle. However, the present disclosure need not be so limited. For example, other types of user devices **105** may include a smartphone, a smartwatch, a tablet computer, a smart speaker, a headset, a television, a thermostat, or the like.

The speech recognition unit **110** includes an utterance detection unit **111** and a grammar **128**. The utterance detection unit **111** can analyze audio data detected by one or more microphones for the occurrence of an utterance **112** by a user, (in this case, the user's name is Bill). The utterance **112** may include a command such as "Turn on AC." Other types of commands may include, for example, "Tune to 99.7," "Increase the volume," "Roll down the window," or the like. In some implementations, the utterance detection unit **111** may be configured to perform additional functionality. For example, the utterance detection unit **111** may be configured to generate a transcription of audio data that the utterance detection unit **111** captures using the one or more microphones. The transcription may include a text version of the audio data that was captured by the one or more microphones in response to Bill's utterance.

5

The speech recognition unit **110** can use the grammar **128** to identify a particular operation that should be performed in response to the uttered command. The grammar **128** may include a set of one or more rules that can be applied to one or more words of the transcription of an utterance to determine the meaning of the utterance. In some implementations, the grammar **128** may use a set of one or more rules to associate one or more particular words with a particular operation. For example, the speech recognition unit **110** can identify a particular operation **128a** based on a comparison of one or more words of the transcription to one or more terminals and non-terminal of a grammar rule.

The speech recognizer **110** can provide data representing the particular operation **128a** to an instruction generation unit **129**. The instruction generation unit **129** can generate a data object **140** (e.g., instructions or commands), based on the received data representing the particular operation **128a**, that, when executed by control system **141**, results in the causing the user device to perform the particular operations.

The control system **141** may include one or more processing units, one or more memory units, and software code that is configured to receive a data object **140**, process the data objection **140**, and instruct one or more components of a user device **105** to perform the particular operation. In some implementations, the control system **141** may include one or more primary processing units, the one or more primary memory units, and software code of the user device **105**. Alternatively, the one or more control units may be a separate subsystem of a user device such as climate control circuitry of a vehicle, a home, or the like. In such implementations, the control system **141** may include one or more processing units, one or more memory units, software code, or a combination thereof that is (or are) not the primary processing unit or the processing memory units. Alternatively, the control system **141** may be only software that is used to start, stop, or both, particular operations that can be performed, or are being performed, respectively, by the user device **105**.

By way of example, the grammar **128** may associate the set of words “turn on AC” with an operation of turning on the air conditioner in a vehicle. Accordingly, in response to the speech recognizer **110** generating a transcription of an uttered voice command that includes the set of words “turn on AC,” the speech recognizer can apply the grammar **128** rules to determine that the vehicle should perform the operation of turning on the air conditioner in the vehicle. The speech recognizer can provide data representing the operation of “turn on the air conditioner” to the instruction generation unit **129**. The instruction generation unit **129** can generate a data object that, when executed by a control system **141**, results in the air conditioner being turned on and used to cool the vehicle. Digital assistant technology associated with the speech recognizer **110** may output a message to Bill through a speaker of the vehicle which recites “Thanks Bill. Turning on AC Now.”

The system **100** of FIG. 1 can perform well offline for a limited set of offline commands that are programmed into the grammar for all users of the device so long as a user such as Bill knows the particular command to initiate performance of the particular operation that Bill wants performed by the user device **105**. However, the system **100** of FIG. 1 fails when a user such as Bill wants to use his own personalized language to perform an operation.

For example, when Bill wants to turn on the air conditioner in his home, Bill may use personalized language to ask another family member in his home to turn on the air conditioner by telling the family member to “freeze me!” As

6

a result, when Bill climbs into his vehicle on a hot day Bill may reflexively say “freeze me!” in an effort to get the air conditioner to turn on and start cooling down the inside of the vehicle. If Bill uttered such an offline command in a vehicle equipped with the system of FIG. 1, then the vehicle’s climate and entertainment system would not respond. Alternatively, vehicle’s climate and entertainment system (or related digital assistant technology) may respond with a message to the effect of “Sorry, I do not understand that command. Please submit another command,” or the like. This is because the grammar of the speech recognizer used by the user device **105** does not include a set of rules for resolving the uttered command of “Freeze Me.” The system of FIG. 2A provides a system that enables the speech recognizer’s **110** grammar **128** to learn Bill’s personalized command of “Freeze Me” so that Bill may utter an offline command “Freeze Me” and have the vehicle turn on the air conditioner in response.

FIG. 2A is a contextual diagram of an example of a system **200A** for learning offline voice commands. The system **200A** includes a user device **105**, a network **150**, and a server **160**. The user device **105** may include an utterance detection unit **111**, one or more logic determination units **113**, **114**, a network interface **115**, a user feedback input unit **119**, a local grammar **128**, an instruction generation unit **129**, and a deletion unit **131**.

The server **160** may include an utterance receiving unit **116**, a command analysis unit **117**, a feedback data receiving unit **120**, one or more logic determination units **123**, **125**, a grammar update unit **124**, and an update delay unit **127**. In this specification, the term “unit” will be used broadly to refer to a software system or subsystem that, when executed by one or more computers, cause the performance of one or more specific functions. Generally, a unit will be implemented as one or more software modules or components installed on the same computer or computers; in other cases, one or more software modules or components may be installed on one or more computers in one or more locations.

With reference to the contextual example of FIG. 2A, a user such as Bill can decide that the temperature is too hot and can utter a command to turn on the air conditioner. For example, a user Bill, speaking in his own vernacular, may utter the command “Freeze Me” **112** in an attempt to turn on the air conditioner. The command “Freeze Me” **112**, at the time of Bill’s utterance of the command “Freeze Me” **112**, was not previously associated with one or more rules of the local grammar **128** of the user device **105**.

One or more microphones of the user device **105** may detect the uttered command “Freeze Me” **112** and generate audio data **112a** based on the detected utterance. The generated audio data **112a** may include an electrical analog signal that is proportional to the acoustic sound wave created by Bill’s voice when Bill utters the command “Freeze Me” **112**. The utterance detection unit **111** can analyze the audio data **112a** generated by the one or more microphones of the user device **105** and detect the occurrence of the command “Freeze Me” **112**. In some implementations, the utterance detection unit **111** can generate a transcription **112b** of the audio data **112a** and perform one or more of the operations below based on the transcription **112b** of the audio data **112a**. Alternatively, in other implementations, the user device **105** may perform one or more of the operations described below on the audio data **112a**, and then provide the audio data **112a** to the server **160** for transcription and further processing by one or more units of the server **160** described herein.

The user device **105** may use a first logic unit **113** to determine whether the server **160** (or another server such as a content server) is required to disambiguate and perform one or more operations in response to the command “Freeze Me” **112**. For example, the first logic unit **113** may be configured to receive the audio data **112a**, the transcription **112b**, or both, and analyze the received audio data **112a**, the received transcription **112b**, or both, to determine whether the server **160** is required to disambiguate and perform one or more operations in response to the command “Freeze Me” **112**.

The user device **105** may determine that the server **160** (or another server such as a content server) is required to execute the command “Freeze Me” **112** for a variety of different reasons. For example, the user device **105** may determine that the server **160** (or another server such as a content server) is required to execute a command that inherently requires data from a remote server such as server **160** (or another server such as a content server) to execute the command. Commands that inherently require data from a remote server such as server **160** (or another server such as a content server) in order to execute include commands such as “What is the weather like in Paris today,” “Order me a large pizza from Tony’s Pizza,” “Tell me the score of the Steeler game,” or the like. Commands such as these inherently require data from a remote server such as server **160** (or another server such as a content server) because each of the respective commands require communication with one or more remote servers to fully execute the received command in a manner that a user such as Bill is likely to find satisfactory (e.g., by retrieving a real-time weather forecast for Paris from a remote server, by placing an order for Pizza with a server of a pizza shop, by retrieving the current score of the Steeler game from a server tracking scores of one or more sporting events).

Another reason the server **160** (or another server such as a content server) may be required to execute the command “Freeze Me” **112** is because the user device **105** may determine that the local grammar **128** of the user device **105** does not include one or more rules that define the operation that should be performed based on the command “Freeze Me.” In such instances, the user device **105** may provide the data representing the uttered command such as a transcription **112b** of the uttered command **112**, audio data **112a** representing the uttered command **112** that can be transcribed by the server **160**, or both, to the server for disambiguation. Alternatively, if the user device **105** determines that the local grammar **128** does define one or more rules for disambiguating the command “Freeze Me” **112**, then the user device **105** may use the local grammar **128** to disambiguate the command “Freeze Me” **112**.

If the uttered command is determined to not need a server **160** (or another server such as a content server) and the local grammar **128** defines one or more rules for disambiguating the command, then the user device can use the local grammar **128** to disambiguate and execute the received command **112**. For example, the user device **105** may perform a search of entries in the grammar **128** based on one or more portions of the transcription **112b** of the command **112** to identify one or more rules that define the operation(s) that should be performed based on the received command “Freeze Me” **112**. In the example of FIG. 2A, the local grammar **128** does not define one or more rules for disambiguating the command “Freeze Me” **112** at the time the command “Freeze Me” **112** is uttered by Bill. As a result, the local grammar **128** needs to be updated using the techniques described herein.

Alternatively, if the user device **105** needs a server **160** (or another server such as a content server) to disambiguate and execute the command **112**, then the user device **105** may use a second logic unit **114** to determine whether the user device **105** has access to a network **150** to communicate with the server **160** (or another server such as a content server). For example, the user device **105** may determine whether the network interface unit **115** of the user device **105** is within range of, and has permission to access, one or more wireless networks of the network **150** such as a WiFi network, a cellular network, or the like.

If the user device **105** determines that it does not have access to the network **150**, the user device can attempt to execute the command “Freeze Me” **112** using the local grammar **128** of the user device **105**. However, if the local grammar **128** does not include any rules for resolving the uttered command of “Freeze Me,” then the user device may provide an output to Bill indicating that the command cannot be executed. Alternatively, if the user device **105** determines that the user device **105** has access to the network **150**, the user device **105** can communicate with the server **160** for assistance related to disambiguating and executing the command “Freeze Me” **112**.

In the example of FIG. 2A, the user device **105** determines that the user device **105** is connected to the network **150** and that the server **160** is required to disambiguate the command “Freeze Me” **112**. In the example of 2A, the server **160** is required to help disambiguate and execute the command “Freeze Me” **112** since the local grammar **128** does not include one or more rules for disambiguating the command “Freeze Me” **112**. The user device **105** may use the network interface unit **115** to transmit the transcription **112b** of the uttered command **112**, the audio data **112a** representing the uttered command **112**, or both, to the server **160** for disambiguation.

With reference to FIG. 2A, the user device **105** is disclosed as including two logic units **113**, **114** for illustrative purposes. However, the present disclosure need not be so limited. For example, a single logic unit may be used to provide the functionality of logic unit **113**, **114** respectively.

The server **160** can receive the transcription **112b** of the uttered command **112**, the audio data **112a** representing the uttered command **112**, or both, using an utterance data receiving unit **116**. The utterance data receiving unit **116** is an interface between the server **160** and the user device **105**. In some implementations, such as when the utterance receiving unit **116** receives audio data **112a** representing the uttered command **112**, the utterance data receiving unit **116** can be configured to generate a transcription of the uttered command **112** based on the audio data **112a**. The utterance data receiving unit **116** may provide the transcription **112b** of the uttered command **112** to the command analysis unit **117**. In other implementations, such as the example of FIG. 2A, the utterance data receiving unit **116** may receive a transcription **112b** of the uttered command **112** and the forward the transcription **112b** to the command analysis unit **117**.

The command analysis unit **117** can be used to disambiguate a command that was not previously defined by the local grammar **130** of the user device. The command analysis unit **117** may disambiguate a command in a number of ways. For example, the command analysis **117** may associate the command “Freeze Me” with turning on the air conditioner based on an analysis of the keywords of the command, based on a user’s response to a request for feedback from Bill regarding the command, based on historical aggregation of Bill’s feedback related to the com-

mand, based on a comparison of the keywords of the command to a grammar maintained by the server 160 that is more extensive than the grammar 130 stored on the user device 105 and has already learned that the command, or a combination thereof.

With respect to the example of FIG. 2A, the command analysis unit 117 may determine, based on an analysis of the keywords in the command “Freeze Me” 112, that the command 112 is associated with a particular operation such as turning on an air conditioner. Based on this determination, the command analysis unit 117 may generate and transmit prompt data 118 for user feedback from Bill that identifies a particular operation. For example, the command analysis unit may generate and transmit prompt data, that when processed for output by the user device 105, asks Bill whether Bill wants to “Turn on the AC?”. The prompt data 118 is received by the network interface unit 115 and output by one or more interface units of the device 105. The interface units may include a speaker, a graphical user interface, or both.

For example, prompt data 118 may be received by the user device 105 and processed to generate the output 118a phrase “Turn on AC?”. In some implementations, the output 118a may be an audio prompt that is output by the one or more speakers of the user device 105. In other implementations, the output may be text rendered on the display of a user interface of the user device 105. In yet other implementations, the output 118a may be provided via a speaker or display of a different user device that is associated with Bill and the user device 105 via a cloud service. The different user device may include, for example, a smartphone, a tablet, a smartwatch, or the like.

Bill may provide feedback 119a in response to the output 118a generated based on the prompt data 118 that indicates whether or not the particular operation identified by the command analysis unit 117 should be associated with the command 112 uttered by Bill. For example, Bill can provide feedback indicating that the particular operation (e.g., turn on air conditioner) identified by the command analysis unit 117 should be associated with the command 112 (e.g., “Freeze Me”) uttered by Bill. Alternatively, Bill can provide feedback indicating that the particular operation (e.g., turn on air conditioner) identified by the command analysis unit 118 should not be associated with the command 112 (e.g., “Freeze Me”). In the example of FIG. 2A, Bill’s feedback 119a includes an utterance that is detected by one or more microphones of the user device 105. Alternatively, or in addition, Bill’s feedback may also include a selection of one or more options of a graphical user interface of the user device 105, an utterance detected by a different user device such as a smartphone, a tablet, a smartwatch, or the like that is associated with Bill and the user device 105 via a cloud service.

In the example of FIG. 2A, Bill may provide feedback to the output 118a prompt “Turn on the AC?” by uttering the word “Yes” 119a. One or more microphones of the user device 105 may detect the Bill’s utterance and generate audio data based on the detected utterance. The generated audio data may include an electrical analog signal that is proportional to the acoustic sound wave created by Bill’s voice when Bill utters the feedback 119 “Yes.” A user feedback input unit can receive the audio data generated based on the uttered feedback 119a, analyze the audio data generated based on the uttered feedback 119a, and detect the occurrence of the feedback 119 phrase “Yes.” In some implementations, the user feedback input unit 119 can generate a transcription 119bn of the audio data generated

based on the uttered feedback 119a and perform one or more of the operations based on the transcription 119b of the audio data generated based on the uttered feedback 119a. The user device 119b may use the network interface unit 115 to transmit the transcription 119b to the server 160 using one or more networks 150. Alternatively, in other implementations, the user device 105 provide the audio data generated based on the uttered feedback 119a to the server 160 via the network 150. In such instances, the server 160 may generate the transcription of the received audio data generated based on the uttered feedback 119a.

A feedback data receiving unit 120 may receive the transcription 119b of the uttered feedback 119a and provide the transcription 119b to the command analysis unit 117. Alternatively, the feedback data receiving unit 120 may receive audio data generated based on the uttered feedback 119a, generate a transcription 119b of the audio data generated based on the uttered feedback 119a, and then provide the transcription 119b of the uttered feedback 119a to the command analysis unit 117.

In response to receiving the transcription 119b of the uttered feedback 119a, the command analysis unit 117 may transmit response data 121 to the user device 105 via the network 150. The user device 105 can receive the response data 121 using the network interface unit 115 and process response data 121 to generate an output 133 confirming the association of the command “Freeze Me” 112 with the particular operation of “turning on the air conditioner” of the environment associated with the user device 105. The output 133 may include output provided via a speaker of the user device 105, output provided via a display of the user device 105, output provided by a speaker of a different user device (e.g., smartphone, tablet, smartwatch, or the like), output provided by a display of a different user device (e.g., a smartphone, tablet, smartwatch, or the like), or a combination thereof. In some implementations, the output 133 may include output that identifies the particular operation that the command analysis unit 117 associated with the command “Freeze Me.” For example, the output 133 may include output from a speaker of the user device 105 that vocalizes output of “Thanks Bill. Turning on the AC Now.”

The command analysis unit 117 can initiate generation of an update to the local grammar 128 of the user device 105. In some implementations, initiating generation of an update to the local grammar 128 of the user device 105 may include, for example, instructing a third logic unit 123 to evaluate data 122 identifying the uttered command 112 to determine whether the uttered command requires the server 160. In some implementations, the determination of whether the uttered command 112 requires the server 160 may be based on the particular operation the command analysis unit 117 associated with the uttered command 112. If the third logic unit 123 determines that the uttered command 112 (or the particular operation the command analysis unit 117 associated therewith) requires the server 160, then the server 160 may cease generation of the update to the local grammar 128. Alternatively, if the third logic unit 123 determines that the uttered command 112 and the particular operation the command analysis unit 117 associated therewith does not require the server 160, then the result of the third logic unit’s 123 processing may provide an indication to the grammar update unit 124 to generate an update for the local grammar 128 of the user device 105.

Generating the update can include generating a plurality of instructions that are configured to, upon execution of the instructions by the user device 105 during processing of the update, create a set of one or more rules in the local grammar

128 that associate the user command (e.g., “Freeze Me”) with the particular operation (e.g., “turn on the air conditioner”) in the local grammar **128** of the user device **105**. In some implementations, generating the local grammar update may also include a generating a data object that, when executed by the user device **105**, causes the user device **105** to initiate performance of the particular operation (e.g., “turn on air conditioner”) responsive to a subsequent utterance of the command **112** (e.g., “Freeze Me.”). The data object may include an XML data object, a binary data object, or the like. In some implementations, generating the local grammar update may also include generating expiration data that includes a maximum period of time that the set of one or more rules associated with the local grammar update will be stored in the local grammar **128** without being used.

Each updating methodology provides particular advantages. For example, generating an update that includes a plurality of instructions to create a set of one or more rules in the local grammar **128** to associate the user command (e.g., “Freeze Me”) with the particular operation (e.g., “turning on the air conditioner”) without the generated data object results in an update package of a smaller size (relative to an update package that includes the generated data object). Accordingly, this update methodology may be used to conserve bandwidth less data would need to be provided to the user device **105** to facilitate updating the local grammar **128**. In addition, update download times may be faster by downloading the plurality of instructions without the generated data object. In addition, this updating methodology also leads to better recognition as there are fewer possible matches. For example, an utterance such as “Freeze We,” “Sneeze Me,” or the like may also be corrected to “Freeze Me” if the words of the utterance do not match words in the local grammar **128**. However, this update methodology will require an instruction generation unit **129** to generate the data object.

Alternatively, different benefits can be achieved by generating the data object and providing the data object with the grammar update for storage on the user device **105** may increase performance of the user device **105**. This is because it saves the user device **105** from having to generate the data object in response the user device **105** identifying the particular operation as being associated subsequent utterances of the command **112**. Because the user device **105** can merely access the data object (as opposed to generating the data object), such embodiments can reduce the latency involved in the user device executing the particular operation responsive to subsequent utterances of the command “Freeze Me.”

Yet different benefits can be achieved by generating and providing expiration data with the local grammar update. For example, the system **200A** can use the expiration data as a mechanism for optimizing use of storage space by the user device **105**. That is, the user device **105** can free-up storage space allocated to store a set of rules associating an uttered command with a particular operation if Bill does not utter the command within the maximum period of time set forth in the expiration data.

The grammar update unit **124** produces the local grammar update **126** and uses a fourth logic unit **125** to determine whether or not the server **160** has access to a network **150** that facilitates communication with the user device **105**. If the fourth logic unit **125** determines that the server **160** does not have access to a network **150** that facilitates communication with the user device **105**, then the fourth logic unit **125** instruct a local grammar update delay unit **127** that the server **160** cannot communicate with the user device **105** at

this time. The local grammar update delay unit **127** can delay subsequent attempts to provide the local grammar update **126** to the user device **105** for a predetermined amount of time. In response to the expiration of the predetermined amount of time, the update delay unit **127** can use the fourth logic unit **125** to determine whether or not the server **160** has access to a network **150** that facilitates communication with the user device **105**.

Alternatively, if the fourth logic unit **125** determines that the server **160** does have access to a network **150** that facilitates communication with the user device **105**, then the server **160** can provide the local grammar update **126** to the user device **105**. The user device **105** can receive the local grammar update via the network interface unit **115**.

The user device **105** may process the instructions included in the local grammar update **126** and update the local grammar **128** by generating a set of one or more rules for storage in the local grammar **128**. The generation and storage of the set of one or more rules in the local grammar **128** enable the local grammar **128** to disambiguate subsequently received commands that correspond one of the commands that were part of the local grammar update **126**. For example, the one or more rules can be applied to one or more words of the transcription of a subsequent utterance to determine the meaning of the subsequent utterance. Determining a meaning of a subsequent utterance may include, for example, a identifying an operation that should be performed in response to the detection of the subsequent utterance. Accordingly, after receiving and processing the local grammar update **126**, the user device **150** can use the local grammar **128** to disambiguate subsequent utterances of the command “Freeze Me” by using one or more rules in the local grammar **128** to associate the “Freeze Me” command with the particular operation of “turning on the air conditioner.”

The instruction generation unit **129** can also be updated based on the local grammar update **126**. In some implementations, for example, the local grammar update **126** may include an action that is to be initiated, performed, or both, by the user device **105** in response to the an uttered command. The action may include a data object that, when executed by the user device **105**, causes the user device **105** to initiate performance of (or to perform) the particular operation (e.g., “turning on air conditioner”) that is associated with an uttered command (e.g., “Freeze Me.”) by the local grammar **128**. The data object may include an XML data object, a binary data object, or the like. If the local grammar update **126** includes the action that that is to be initiated, performed, or both, by the user device **105** in response to an uttered command, the action can be stored in the instruction generation unit **129**. Once stored in the instruction generation unit **129**, the action (e.g., data object that instructs the user device **105** to turn on the air conditioner) may be retrieved upon after receipt of the utterance of a corresponding command (e.g., “Freeze Me”) based on data provided by the local grammar **128** that identifies a particular operation associated with the utterance. Alternatively, the action (e.g., data object that instructs the user device **105** to turn on the air conditioner) may be generated by the instruction generation unit **129** after receipt of the utterance of a corresponding comment (e.g., “Freeze Me”) based on data provided by the local grammar **128** that identifies a particular operation associated with the utterance.

Other units may also be updated based on the local grammar update **126**. For example, the deletion unit **131** may be updated to monitoring the usage of one or more sets

13

of rules that are included in the local grammar **128**. In some implementations, a set of grammar rules may be associated with a maximum period of time for which the set of grammar rules may be idle (e.g., not used). The deletion unit **131** may determine whether or not a command that is associated with a set of grammar rules has been uttered within the maximum period of time. In response to determining that the command has not been uttered, or otherwise input, within the maximum period of time, the deletion unit **131** may instruct **132** the local grammar **128** to delete the set of rules from the local grammar **128**. This feature provides the benefit of optimizing the use of local user device storage, which may be finite.

Though deletion unit **131** has been described as removing a set of one or more rules from the local grammar **128** based on a determination that a command associated with the set of one or more rules has not been uttered, or otherwise input, within a maximum period of time, the present disclosure need not be so limited. Instead, the deletion unit **128** may detect the occurrence of one or more other triggers whose detection initiates removal (e.g., deletion) of a set of one or more rules associated with a particular command in the local grammar **128**. Such triggers may include changes in the state of the user device **105** such as removal of a user account, selection of a different language for the user device **105**, a deletion request received from a user (e.g., Bill), a request to reset the device received from a user (e.g., Bill), or the like.

In the example of FIG. 2B, the user device that detected Bill's utterance of a previously unknown command (e.g., "Freeze Me") and the user device that received and processed the local grammar update **126** were the same user device—i.e., user device **105**. However, the present disclosure need not be so limited. Instead, for example, multiple user devices may be associated with Bill via one or more cloud services. In such instances, the server **160** can transmit the local grammar update such as local grammar update **126** to each of the other devices associated with Bill via the one or more cloud services. For example, the local grammar update can also be transmitted to Bill's smartphone, smartwatch, laptop computer, tablet computer, desktop computer, smart television, set top box, smart speaker, digital home assistant, or the like. In this way, the local grammar of each of Bill's user devices can be updated to respond to commands in Bill's own vernacular. Once updated, these user devices can be used to respond to commands in Bill's own vernacular even when offline.

FIG. 2B is a contextual diagram of the system of FIG. 2A that has learned an offline voice command. In general, the system **200B** includes a user device **111**, a first logic unit **113**, a second logic unit **114**, a grammar **128**, an instruction generation unit **111**, a deletion unit **131**, a feedback input unit **119**, and a network interface unit **115**.

In more detail, the system **200B** is the same as the system **200A** except the system **200B** is shown after the local grammar of **128** has received and processed the local grammar update **126**. Processing of the local grammar update **126** enabled the local grammar **128** to learn a command (e.g., "Freeze Me") that was not associated within one or more rules in the local grammar **128** prior to receiving and processing the local grammar update **126**. In addition, the user device **105** of FIG. 2B does not have access to the network **150**. Because the user device **105** of FIG. 2B does not have access to the network **150**, the user device **105** of FIG. 2B cannot communicate with the server **160** to avail itself of the command disambiguation services provided by the server **160**.

14

With reference to FIG. 2B, the user device **105** may detect utterance of a command **112** that is uttered in Bill's own vernacular such as the command **112** "Freeze Me" that was unknown to the local grammar **128** prior to the user's device's **105** receipt and processing of the local grammar update **126**.

The user device **105** may use one or more microphones to capture audio data **112a** that is provided to the utterance detection unit **111**. The utterance detection unit **111** may generate a transcription **112b** of the audio data **112a**. The user device **105** may use the first logic unit **113** to determine that the server is not needed to execute the command represented by the transcription **112b** because the command does not require data from the server **160** and the local grammar **128** that was updated using the local grammar update **126** now includes the command **112** represented by the transcription **112b**. The second logic unit **114** may determine that there is no access to the network **150**.

The local grammar **128** may determine, based on the received a transcription **112b** and a set of one or more rules stored in the local grammar that were received as part of the local grammar update **126**, that the command **112** (e.g., "Freeze Me") is associated with a particular operation (e.g., turning on the air conditioner). The local grammar may transmit data **228a** instruction generation unit **129** that identifies the particular operation (e.g., turning on the air conditioner) that is associated with the command **112**.

The instruction generation unit **129** may obtain an action **130**, based on the particular operation identified in the data **228a** received from the local grammar **128**, that when, executed, initiates performance of a particular operation (e.g., turning on the air conditioner) that should be performed in response to the command **112**. In some implementations, the action **130** may have been previously received and stored by the instructions generation unit **129** as part of the local grammar update **126**. In such implementations, the latency for executing the particular operation in response to the command **112** may be reduced by storing the action **130** on the user device **150**. Alternatively, obtaining the action **130** may include, for example, generating the action **130**.

The action **130** may include a data object that, when executed by the user device **105**, causes the user device **105** to initiate performance of the particular operation (e.g., "turning on air conditioner") responsive to a subsequent utterance of the command **112** (e.g., "Freeze Me."). The data object may include an XML data object, a binary data object, or the like. The action **130** may be processed by the user device **105** and trigger the performance of the particular operation (e.g., turning on the air conditioner).

FIG. 3 is a flowchart of an example of a process **300** for generating a personalized grammar update for a user device. For convenience, the process **300** is described below as being performed by a system such as the system **200A** depicted in FIG. 2A.

The system may use a server to receive **310** data representing at least a portion of an utterance from a first user device. The data representing the portion of the utterance may include a transcription of one or more portions of the utterance. Alternatively, the data representing the portion of the utterance may include audio data representing one or more portions of the utterance. The utterance may include a command in the unique vernacular of a user for a known operation. For example, the utterance may include a command to "Hurt My Ears." In this example, the user may utter the command "Hurt My Ears" as a request to perform the

15

operation of “increasing the volume of audio (or video) playback by the first user device.”

The system may determine **320**, based on the data representing the utterance, an operation that can be performed responsive to the utterance. In some implementations, the operation may be an operation that can be performed without requesting server information. For example, the system may determine that the operation of “increasing the volume of audio (or video) playback by a device” is an operation that can be performed without requesting server information.

The system may generate **330** an update to the grammar based on the (I) data representing at least a portion of the utterance and (II) the operation. Generating an update to the grammar may include generating a plurality of instructions that are configured to, upon execution of the instructions by the first user device during processing of the update, create a set of one or more rules in a grammar of the first user device that associates the user command (e.g., “Hurt My Ears”) with the particular operation (e.g., “increasing the volume of audio (or video) playback by the first”).

In some implementations, generating the update to the grammar may also include generating a data object that, when executed by the first user device, causes the first user device to initiate performance of the particular operation (e.g., “increasing the volume of audio (or video) playback by the first user device”) responsive to a subsequent utterance of the command. The data object may include an XML data object, a binary data object, or the like. In some implementations, generating the grammar update may also include generating expiration data that includes a maximum period of time that the set of one or more rules associated with the grammar update will be stored in the grammar without being used.

The system may determine **340** whether the first device and the first device and the server are connected by one or more networks. In response to determining that the first user device and the server are connected by one or more networks, then the system may provide **350** the grammar update to the first user device.

Alternatively, in response to determining that the first user device and the server are not connected by one or more networks, the system may proceed to stage **342** and delay **342** the grammar update until the first user device and the server are connected by one or more networks. In such instances, the system may continue to stage **350** and provide the grammar update to the first user device after determining that a connection has been established between the first user device and the server.

FIG. 4 is a flowchart of an example of a process **400** for personalizing a grammar of a user device. For convenience, the process **400** is described below as being performed by a system such as the system **200A** depicted in FIG. 2A.

The system can detect **410**, using a first user device, an utterance of one or more terms when a user is connected to a server by a network. For example, the first user device may use one or more microphones to capture audio data that represents the utterance. By way of example, the utterance may include a command in the user’s own vernacular to “Hurt My Ears.”

The system can provide **420**, using the first device, data representing the utterance to the server using the network. For example, the system may provide the captured audio data describing all, or a portion of, the utterance to the server. Alternatively, or in addition, the system may provide a transcription of all, or a portion of the utterance to the server. By way of example, the first device may provide a

16

transcription of audio data that represents the command “Hurt My Ears” to the server.

The system can receive **430**, by the user device and from the server, an update to the grammar of the user device based on the data representing the utterance. Receiving the update to the grammar of the user device may include obtaining a copy of update to the grammar via a network. Alternatively, or in addition, receiving the update to the grammar may also include installing the update to the grammar. The update to the grammar of the user device may include a plurality of instructions that are configured to, upon execution of the instructions by the first user device during processing of the update, create a set of one or more rules in a grammar of the first user device that associates the user command (e.g., “Hurt My Ears”) with the particular operation (e.g., “increasing the volume of audio (or video) playback by the first”).

In some implementations, the update to the grammar may also include a data object that, when executed by the first user device, causes the first user device to initiate performance of the particular operation (e.g., “increasing the volume of audio (or video) playback by the first user device”) responsive to a subsequent utterance of the command. The data object may include an XML data object, a binary data object, or the like. In some implementations, the grammar update may also include expiration data that includes a maximum period of time that the set of one or more rules associated with the grammar update will be stored in the grammar without being used.

The system can detect **440** a subsequent utterance that includes the one or more terms when the first user device is not connected to the server by a network. For example, the first user device may use one or more microphones to capture audio data that represents the subsequent utterance. The subsequent utterance may include an utterance that was received after the receipt and processing of the update to the grammar in stage **430**. For example, the subsequent utterance may include the command “Hurt My Ears” that was uttered by the user at some point in time after the update to the grammar was received in stage **430**.

The system can determine **450**, using the first user device, that the subsequent utterance is associated with an operation that can be performed without requesting information from the server. For example, the first user device can evaluate data representing the utterance to determine whether information from a server is required to disambiguate the utterance and perform the operation that is responsive to the utterance.

The system can obtain **460** a data object that can be used to trigger performance of the operation based on (I) the one or more terms of the subsequent utterance and (II) the updated grammar. The data object may include data that, when executed by the first user device, causes the first user device to initiate performance of an operation (e.g., “increasing the volume of audio (or video) playback by the first user device”) that was associated with the utterance in the updated grammar. The data object may include an XML data object, a binary data object, or the like. The system can initiate **470**, using the user device, performance of the operation without requesting information from the server by processing the data object.

Embodiments of the subject matter, the functional operations and the processes described in this specification can be implemented in digital electronic circuitry, in tangibly-embodied computer software or firmware, in computer hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or

more of them. Embodiments of the subject matter described in this specification can be implemented as one or more computer programs, i.e., one or more modules of computer program instructions encoded on a tangible nonvolatile program carrier for execution by, or to control the operation of, data processing apparatus. Alternatively, or in addition, the program instructions can be encoded on an artificially generated propagated signal, e.g., a machine-generated electrical, optical, or electromagnetic signal that is generated to encode information for transmission to suitable receiver apparatus for execution by a data processing apparatus. The computer storage medium can be a machine-readable storage device, a machine-readable storage substrate, a random or serial access memory device, or a combination of one or more of them.

The term “data processing apparatus” encompasses all kinds of apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, or multiple processors or computers. The apparatus can include special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit). The apparatus can also include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, or a combination of one or more of them.

A computer program (which may also be referred to or described as a program, software, a software application, a module, a software module, a script, or code) can be written in any form of programming language, including compiled or interpreted languages, or declarative or procedural languages, and it can be deployed in any form, including as a standalone program or as a module, component, subroutine, or other unit suitable for use in a computing environment. A computer program may, but need not, correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, subprograms, or portions of code). A computer program can be deployed to be executed on one computer or on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communication network.

The processes and logic flows described in this specification can be performed by one or more programmable computers executing one or more computer programs to perform functions by operating on input data and generating output. The processes and logic flows can also be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit).

Computers suitable for the execution of a computer program include, by way of example, can be based on general or special purpose microprocessors or both, or any other kind of central processing unit. Generally, a central processing unit will receive instructions and data from a read-only memory or a random access memory or both. The essential elements of a computer are a central processing unit for performing or executing instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., mag-

netic, magneto optical disks, or optical disks. However, a computer need not have such devices. Moreover, a computer can be embedded in another device, e.g., a mobile telephone, a personal digital assistant (PDA), a mobile audio or video player, a game console, a Global Positioning System (GPS) receiver, or a portable storage device (e.g., a universal serial bus (USB) flash drive), to name just a few.

Computer readable media suitable for storing computer program instructions and data include all forms of nonvolatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto optical disks; and CD-ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

To provide for interaction with a user, embodiments of the subject matter described in this specification can be implemented on a computer having a display device, e.g., a CRT (cathode ray tube), LCD (liquid crystal display), LED (light-emitting diode), or OLED (organic light-emitting diode) monitor, for displaying information to the user and a keyboard and a pointing device, e.g., a mouse, a trackball, or a touchscreen, by which the user can provide input to the computer. Other kinds of devices can be used to provide for interaction with a user as well; for example, feedback provided to the user can be any form of sensory feedback, e.g., visual feedback, auditory feedback, or tactile feedback; and input from the user can be received in any form, including acoustic, speech, or tactile input. In addition, a computer can interact with a user by sending documents to and receiving documents from a device that is used by the user; for example, by sending web pages to a web browser on a user's user device in response to requests received from the web browser.

Embodiments of the subject matter described in this specification can be implemented in a computing system that includes a back end component, e.g., as a data server, or that includes a middleware component, e.g., an application server, or that includes a front end component, e.g., a client computer having a graphical user interface or a Web browser through which a user can interact with an implementation of the subject matter described in this specification, or any combination of one or more such back end, middleware, or front end components. The components of the system can be interconnected by any form or medium of digital data communication, e.g., a communication network. Examples of communication networks include a local area network (“LAN”) and a wide area network (“WAN”), e.g., the Internet.

The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other.

While this specification contains many specific implementation details, these should not be construed as limitations on the scope of what may be claimed, but rather as descriptions of features that may be specific to particular embodiments. Certain features that are described in this specification in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable subcombination. Moreover, although features may be described

19

above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a subcombination or variation of a subcombination.

Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above should not be understood as requiring such separation in all embodiments, and it should be understood that the described program components and systems can generally be integrated together in a single software product or packaged into multiple software products.

Particular embodiments of the subject matter have been described. Other embodiments are within the scope of the following claims. For example, the actions recited in the claims can be performed in a different order and still achieve desirable results. As one example, the processes depicted in the accompanying figures do not necessarily require the particular order shown, or sequential order, to achieve desirable results. In certain implementations, multitasking and parallel processing may be advantageous. Other steps or stages may be provided, or steps or stages may be eliminated, from the described processes. Accordingly, other implementations are within the scope of the following claims.

The invention claimed is:

1. A method comprising:

receiving a transcription of an utterance that was detected via a first user interface input, wherein the transcription of the utterance includes one or more terms corresponding to the utterance;

identifying, based on the terms of the transcription of the utterance, an operation to be performed responsive to the utterance, wherein the operation to be performed responsive to the utterance comprises changing a state of a user device comprising the first user interface;

determining, whether the operation can be performed at the user device responsive to the utterance and without requesting information from a server, wherein determining that the operation can be performed by the user device responsive to the utterance and without requesting information from the server comprises:

determining that the user device does not require data from the server to change the state responsive to the utterance, and

determining that the user device includes one or more rules associated with changing the state responsive to the utterance; and

based on determining, that the operation can be performed by the user device responsive to the utterance and without requesting information from the server:

causing, based on the one or more terms of the transcription corresponding to the utterance and one or more of the rules associated with the operation to be performed responsive to the utterance, an update to a grammar stored locally on the user device to be provided,

wherein, when it is determined, that the operation cannot be performed by the user device responsive to the utterance and without requesting informa-

20

tion from the server, no update is provided to the grammar stored locally on the user device;

subsequent to, and in response to, causing the update to the grammar stored locally on the user device to be provided:

identifying a user profile that is associated with the user device, wherein identifying the user profile that is associated with the user device comprises identifying the user profile based on the user profile being associated with the user device and the user device being the user device to which the update to the grammar is provided;

identifying, based on the user profile, one or more other user devices that are registered to a user of the user device, wherein identifying the one or more other user devices based on the user profile is in response to the user profile being identified as associated with the user device;

determining, whether the operation can be performed by one or more of the other user devices responsive to the utterance and without requesting information from the server; and

causing, based on the identification of an other user device of the other user devices that are registered to the user of the user device, based on determining the operation can be performed by the other user device responsive to the utterance and without requesting information from the server, and based on the one or more terms, an update to an additional grammar stored locally on the other user device,

wherein, when it is determined, that the operation cannot be performed by the other user device responsive to the utterance and without requesting information from the server, no update is provided to the additional grammar stored locally on the other user device.

2. The method of claim 1, wherein causing the update to the grammar stored locally on the user device to be provided includes causing provision of the one or more terms to the user device.

3. The method of claim 2, further comprising:

associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent receiving a subsequent utterance of the one or more terms by the user to perform the operation.

4. The method of claim 1, further comprising:

causing generation, by the server, of a data object that, when executed, causes the user device to initiate performance of the operation responsive to the utterance.

5. The method of claim 4, wherein causing an update to the grammar stored locally on the user device to be provided further comprises causing provision of the generated data object to the user device.

6. The method of claim 1, wherein the update to the grammar enables the user device to execute a data object that initiates performance of the operation in response to subsequently receiving the utterance, and without requesting information from the server.

7. The method of claim 1, wherein the other user devices comprise one of: a smartphone, a smartwatch, a laptop computer, a tablet computer, a desktop computer, a smart television, a set top box, a smart speaker, and/or a digital home assistant.

21

8. A system comprising:
 one or more computers and one or more storage devices
 storing instructions that are operable, when executed by
 the one or more computers, to cause the one or more
 computers to perform operations comprising:
 receiving a transcription of an utterance that was
 detected via a first user interface input, wherein the
 transcription of the utterance includes one or more
 terms corresponding to the utterance;
 identifying, based on the terms of the transcription of
 the utterance, an operation to be performed respon-
 sive to the utterance, wherein the operation to be
 performed responsive to the utterance comprises
 changing a state of a user device comprising the first
 user interface;
 determining, whether the operation can be performed at
 the user device responsive to the utterance and
 without requesting information from a server,
 wherein determining that the operation can be per-
 formed by the user device responsive to the utterance
 and without requesting information from the server
 comprises:
 determining that the user device does not require
 data from the server to change the state responsive
 to the utterance, and
 determining that the user device includes one or
 more rules associated with changing the state
 responsive to the utterance; and
 based on determining, that the operation can be per-
 formed by the user device responsive to the utterance
 and without requesting information from the server:
 causing, based on the one or more terms of the
 transcription corresponding to the utterance and
 one or more of the rules associated with the
 operation to be performed responsive to the utter-
 ance, an update to a grammar stored locally on the
 user device to be provided,
 wherein, when it is determined, that the operation
 cannot be performed by the user device respon-
 sive to the utterance and without requesting
 information from the server, no update is pro-
 vided to the grammar stored locally on the user
 device;
 subsequent to, and in response to, causing the update
 to the grammar stored locally on the user device to
 be provided:
 identifying a user profile that is associated with the
 user device, wherein identifying the user profile
 that is associated with the user device comprises
 identifying the user profile based on the user
 profile being associated with the user device
 and the user device being the user device to
 which the update to the grammar is provided;
 identifying, based on the user profile, one or more
 other user devices that are registered to a user of
 the user device, wherein identifying the one or
 more other user devices based on the user
 profile is in response to the user profile being
 identified as associated with the user device;
 determining, whether the operation can be per-
 formed by one or more of the other user devices
 responsive to the utterance and without request-
 ing information from the server; and
 causing, based on the identification of an other
 user device of the other user devices that are
 registered to the user of the user device, based
 on determining the operation can be performed

22

by the other user device responsive to the
 utterance and without requesting information
 from the server, and based on the one or more
 terms, an update to an additional grammar
 stored locally on the other user device,
 wherein, when it is determined, that the opera-
 tion cannot be performed by the other user
 device responsive to the utterance and without
 requesting information from the server, no
 update is provided to the additional grammar
 stored locally on the other user device.

9. The system of claim 8, wherein causing the update to
 the grammar stored locally on the user device to be provided
 includes causing provision of the one or more terms to the
 user device.

10. The system of claim 9, the operations further com-
 prising:

associating the one or more terms with a predetermined
 time period, wherein the predetermined time period
 indicates a maximum period of time the one or more
 terms will be included in the grammar absent receiving
 a subsequent utterance of the one or more terms by the
 user to perform the operation.

11. The system of claim 8, the operations further com-
 prising:

causing generation, by the server, of a data object that,
 when executed, causes the user device to initiate per-
 formance of the operation responsive to the utterance.

12. The system of claim 11, wherein causing an update to
 the grammar stored locally on the user device to be provided
 further comprises causing provision of the generated data
 object to the user device.

13. The system of claim 8, wherein the update to the
 grammar enables the user device to execute a data object that
 initiates performance of the operation in response to subse-
 quently receiving the utterance, and without requesting
 information from the server.

14. The system of claim 8, wherein the operations are
 locally performed at the user device and/or the other user
 devices.

15. The system of claim 8, wherein the operations are
 remotely performed at the server.

16. The system of claim 8, wherein a first set of the
 operations are locally performed at the user device and/or
 the other user devices, and a second set of the operations are
 remotely performed at a server.

17. A non-transitory computer-readable medium storing
 software comprising instructions executable by one or more
 computers which, upon such execution, cause the one or
 more computers to perform operations comprising:

receiving a transcription of an utterance that was detected
 via a first user interface input, wherein the transcription
 of the utterance includes one or more terms correspond-
 ing to the utterance;

identifying, based on the terms of the transcription of the
 utterance, an operation to be performed responsive to
 the utterance, wherein the operation to be performed
 responsive to the utterance comprises changing a state
 of a user device comprising the first user interface;

determining, whether the operation can be performed at
 the user device responsive to the utterance and without
 requesting information from a server, wherein deter-
 mining that the operation can be performed by the user
 device responsive to the utterance and without request-
 ing information from the server comprises:

23

determining that the user device does not require data from the server to change the state responsive to the utterance, and
determining that the user device includes one or more rules associated with changing the state responsive to the utterance; and
based on determining, that the operation can be performed by the user device responsive to the utterance and without requesting information from the server:
causing, based on the one or more terms of the transcription corresponding to the utterance and one or more of the rules associated with the operation to be performed responsive to the utterance, an update to a grammar stored locally on the user device to be provided,
wherein, when it is determined, that the operation cannot be performed by the user device responsive to the utterance and without requesting information from the server, no update is provided to the grammar stored locally on the user device;
subsequent to, and in response to, causing the update to the grammar stored locally on the user device to be provided:
identifying a user profile that is associated with the user device, wherein identifying the user profile that is associated with the user device comprises identifying the user profile based on the user profile being associated with the user device and the user device being the user device to which the update to the grammar is provided;
identifying, based on the user profile, one or more other user devices that are registered to a user of the user device, wherein identifying the one or more other user devices based on the user profile is in response to the user profile being identified as associated with the user device;

24

determining, whether the operation can be performed by one or more of the other user devices responsive to the utterance and without requesting information from the server; and
causing, based on the identification of an other user device of the other user devices that are registered to the user of the user device, based on determining the operation can be performed by the other user device responsive to the utterance and without requesting information from the server, and based on the one or more terms, an update to an additional grammar stored locally on the other user device,
wherein, when it is determined, that the operation cannot be performed by the other user device responsive to the utterance and without requesting information from the server, no update is provided to the additional grammar stored locally on the other user device.

18. The non-transitory computer-readable medium of claim 17, wherein causing the update to the grammar stored locally on the user device to be provided includes causing provision of the one or more terms to the user device.

19. The non-transitory computer-readable medium of claim 18, the operations further comprising:
associating the one or more terms with a predetermined time period, wherein the predetermined time period indicates a maximum period of time the one or more terms will be included in the grammar absent receiving a subsequent utterance of the one or more terms by the user to perform the operation.

20. The non-transitory computer-readable medium of claim 17, the operations further comprising:
causing generation, by the server, of a data object that, when executed, causes the user device to initiate performance of the operation responsive to the utterance.

* * * * *