

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388679
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Gupta; Anshuman et al.

Full mesh proxyless connectivity between networks

Abstract

The present disclosure provides for full mesh connectivity between all endpoints in a VPC, including virtual machines, load balancers, routers, interconnects, virtual private networks, etc. Connectivity may be extended to on-premises devices, such as those connected via VPN and Interconnects. The connectivity is high performance, reliable, and secure.

Inventors: Gupta; Anshuman (Sunnyvale, CA), Radoslavov; Pavlin (Cupertino, CA), Sattainathan; Kannan (Fremont, CA), Kumar; Alok (Fremont, CA), Richter; Yossi (San Jose, CA)

Applicant: Google LLC (Mountain View, CA)

Family ID: 84331859

Assignee: Google LLC (Mountain View, CA)

Appl. No.: 17/859558

Filed: July 07, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230379191 A1	Nov. 23, 2023

Related U.S. Application Data

us-provisional-application US 63344842 20220523

Publication Classification

Int. Cl.: H04L12/66 (20060101); H04L45/74 (20220101); H04L67/10 (20220101)

U.S. Cl.:

CPC **H04L12/66** (20130101); **H04L45/74** (20130101); **H04L67/10** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10193968	12/2018	Alpert et al.	N/A	N/A
10999163	12/2020	Abraham	N/A	N/A
2015/0381493	12/2014	Bansal	370/392	H04L 45/741
2018/0069754	12/2017	Dasu et al.	N/A	N/A
2018/0262387	12/2017	Mathew et al.	N/A	N/A
2018/0307522	12/2017	Wu et al.	N/A	N/A
2019/0319847	12/2018	Nahar	N/A	G06F 9/45558
2020/0067734	12/2019	Hira et al.	N/A	N/A
2020/0213225	12/2019	Han	N/A	H04L 47/829
2020/0382345	12/2019	Zhao et al.	N/A	N/A
2024/0223451	12/2023	Almadani	N/A	H04L 41/122

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2963866	12/2015	EP	N/A

OTHER PUBLICATIONS

Extended European Search Report for European Patent Application No. 22204203.8 dated May 15, 2023. 9 pages. cited by applicant

Primary Examiner: Abaza; Ayman A

Assistant Examiner: Santarisi; Abdul Aziz

Attorney, Agent or Firm: Lerner David LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application claims the benefit of the filing date of U.S. Provisional Patent Application No. 63/344,842 filed May 23, 2022, the disclosure of which is hereby incorporated herein by reference.

BACKGROUND

(1) Virtual Private Cloud (VPC) provides both connectivity as well as an administrative boundary in public clouds. However, sometimes a workload in a VPC needs to connect to resources in another VPC in a different administrative domain.

BRIEF SUMMARY

- (2) The present disclosure provides for full mesh connectivity between all endpoints in a VPC. Such endpoints include, for example, virtual machines, load balancers, routers, interconnects, virtual private networks, etc. Connectivity may be extended to on-premises devices, such as those connected via VPN and Interconnects. The connectivity is high performance, reliable, and secure.
- (3) One aspect of the disclosure provides a system, comprising a software-defined networking (SDN) controller, an on-host controller, and a programmable packet switcher, wherein the SDN controller, on-host controller, and programmable packet switcher are communicatively coupled within a first network, wherein the on-host controller is programmed to receive limited connectivity information from a target network, and use the limited connectivity information to establish a proxyless peering connection between the source network and the target network. The system may include a first virtual machine within the first network. The first virtual machine and the on-host controller may reside on a same host machine. In some examples, the first network is a virtual private cloud (VPC).
- (4) The limited connectivity information may comprise a destination address in the target network. It may further comprise forwarding information.
- (5) According to some examples, the programmable packet switcher delivers packets from a first virtual machine in the source network to a second virtual machine in the target network using the limited connectivity information.
- (6) According to some examples, the source network includes a plurality of endpoints, and each endpoint in the source network is adapted to communicate directly with each endpoint in the target network using the proxyless peering connection.
- (7) According to some examples, the system may further comprise a plurality of SDN controllers in a fully sharded control plane, each of the plurality of SDN controller responsible for a subset of virtual private cloud (VPC) networks in a peering group, each VPC network having an associated on-host controller, wherein each of the plurality of SDN controllers is coupled to each of the associated on-host controllers.
- (8) According to some examples, the system may further comprise a scalable gateway between the programmable packet switch of the first network and a second programmable packet switcher of the second network.
- (9) Another aspect of the disclosure provides a method of establishing a peering connection between a first network and a second network. The method may include receiving, at an on-host controller in the first network, programming information from a first controller within the first network, receiving, at the on-host controller in the first network, connectivity information from a second controller within the second network, establishing a proxyless peering connection between the first network and the second network using at least the received connectivity information, and providing, to a packet switcher within the first network, packet information for direct communication between a first endpoint in the first network and a second endpoint in the second network. The first endpoint may be a virtual machine. Receiving the connectivity information from the second controller may include using the programming information from the first controller to subscribe to the second controller. The connectivity information may include a destination address in the second network and forwarding information.
- (10) According to some examples, the method may further include delivering, by the programmable packet switcher, packets from a first virtual machine in the first network to a second virtual machine in the second network using the connectivity information.
- (11) Yet another aspect of the disclosure provides a software-defined networking (SDN) controller executable on one or more processors for performing a method of establishing a peering connection between a first network that includes the SDN controller and a second network. The SDN controller may be configured to provide programming information to an on-host controller in the first network, the programming information identifying the second network and causing the on-host controller to subscribe to a second SDN controller in the second network. According to some

examples, the SDN controller includes a peering multiplexer configured to receive information from a plurality of resources and publish routing information for communicating with the resources. According to some examples, the SDN controller includes a peering manager maintaining one or more export policies for the first network, wherein the peering multiplexer performs multiplexing on the resources based on the export policies. According to some examples, the SDN controller includes a routing manager configured to receive route information from one or more other SDN controllers. According to some examples, the SDN controller is configured to identify route conflicts based on the route information received and to resolve the conflict using a pre-defined set of rules.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a pictorial diagram illustrating connectivity between two VPCs, according to aspects of the disclosure.
- (2) FIG. 2 is a block diagram illustrating cross-VPC connectivity in a software-defined network (SDN) architecture, according to aspects of the disclosure.
- (3) FIG. 3 is a block diagram illustrating example details of cross-VPC connectivity according to aspects of the disclosure.
- (4) FIG. 4 is a block diagram illustrating components of an on-host controller used for establishing peering, according to aspects of the disclosure.
- (5) FIG. 5 is a block diagram illustrating an example hardware environment according to aspects of the disclosure.
- (6) FIG. 6 is a block diagram illustrating connection between a peering group and a SDN data plane via a fully sharded SDN control plane, according to aspects of the disclosure.
- (7) FIG. 7 is a block diagram illustrating route conflict resolution across peered VPCs, according to aspects of the disclosure.
- (8) FIG. 8 is a block diagram illustrating scalable cross-VPC connectivity via gateways, according to aspects of the disclosure.
- (9) FIG. 9 is a block diagram illustrating cross-network multiplexing according to aspects of the disclosure.
- (10) FIG. 10 is a flow diagram illustrating an example method according to aspects of the disclosure.

DETAILED DESCRIPTION

- (11) FIG. 1 illustrates an example of full mesh proxyless connectivity between two VPCs **110**, **120**. As shown, each VPC **110**, **120** includes a plurality of different endpoints, such as virtual machine **112**, virtual private network (VPN) connection **114**, and internal load balancing (ILB) node **116** in first VPC **110**, and virtual machine **122**, interconnect **124**, and ILB **126** in second VPC **120**. While these are some examples of the types of endpoints, it should be understood that other types of endpoints may also be included in the VPCs. One example of another type of endpoint includes routers. While several endpoints are illustrated in each VPC **110**, **120**, it should be understood that any number of endpoints may be included in each VPC **110**, **120**, and that the number of computing devices in a first VPC may differ from a number of computing devices in a second VPC. Moreover, it should be understood that the number of computing devices in each VPC may vary over time, for example, as hardware is removed, replaced, upgraded, or expanded. Each of the endpoints in the first VPC **110** may be communicatively coupled with each of the endpoints in the second VPC **120**.
- (12) The virtual machines **112**, **122** may each represent one or more virtual machines running on one or more host machines. Each of the virtual machines **112**, **122** may run an operating system and applications. While only a few virtual machines **112**, **122** are illustrated, it should be understood

that any number of virtual machines may be supported by any number of host computing devices.

(13) The ILBs **116**, **126** may include either or both of network load balancers, application load balancers, or other types of load balancers. According to some examples, the ILBs **116**, **126** may distribute jobs or tasks among the virtual machines **112**, **122**. For example, the computing devices may have different capacities, such as different processing speeds, workloads, etc. In some examples, as commands are received to execute new tasks, the ILBs **116**, **126** may determine which virtual machine is responsible for handling the new task. Moreover, the ILBs **116**, **126** may monitor computing capacity of each of the virtual machines **112**, **122** and make adjustments as needed. For example, if one virtual machine fails, the ILB **116**, **126** responsible for load balancing for that virtual machine may move its processes over to one or more other virtual machines. In other examples, if one virtual machine is near capacity while another has availability, the ILB **116**, **126** may move jobs, or may only send new jobs to the virtual machine with availability.

(14) The VPN connection **114** may establish a connection between the VPC **110** and one or more on-premise devices **119**. For example, the on-premise devices **119** may include routers or other devices within an organization. The VPN connection **114** establishes a protected network connection between the on-premise devices **119** and the VPC **110**, such as by encrypting internet traffic.

(15) The interconnect **124** may establish a connection between the VPC **120** and the on-premise devices **129**. For example, the interconnect **124** may establish a direct physical connection between an on-premises network and the VPC **120**, such as through a supported service provider.

(16) The on-premise devices **129** may be the same or different than the on-premise devices **119**. The on-premise devices **119** may belong to a same or different organization than the on-premise devices **129**. The VPC peering between the VPC **110** and the VPC **120** may provide for communication of information between the VPCs **110**, **120**. By way of example only, the first VPC **110** may be part of a first organization which is a company that offers a database service, whereas the second VPC **120** may be part of a second organization which is a separate company with a web frontend that uses the database service. As another example, the first VPC **110** may be a first department, such as a marketing department, within an organization, and the second VPC **120** may be a second department within the same organization, such as a sales department. While these are merely two examples of VPCs for which full mesh connectivity may be beneficial, it should be understood that the description herein may apply to any of a wide variety of other contexts.

(17) While only two on-premise networks and VPCs are shown in FIG. 1, it should be understood that multiple on-premise networks and VPCs may connect virtually. Conventionally, a view of the resources in another VPC is provided by copying the resources. This can be costly in terms of time, bandwidth consumption, memory consumption, etc. According to the present disclosure, such copying may be avoided by fetching only the information needed to establish a peering connection with the other VPC.

(18) FIG. 2 illustrates an example of cross-VPC connectivity in a software-defined network (SDN) architecture. The architecture has several layers of control. In a first layer, each SDN controller **232**, **252** is responsible for one of two or more networks **230**, **250**. In a second layer, an on-host controllers (OHC) **234**, **254** in each network **230**, **250** is used to obtain information for peering with the other network. When establishing a peering relationship, the OHC **234**, **254** of the network **230**, **250** seeks only a limited amount of information needed to establish peering, as opposed to copying all information. For example, if VM_A **238** in first network **230** seeks to peer with VM_B in second network **250**, the OHC **234** in first network **230** obtains only enough information from SDN controller **252** in the second network **250** to allow the VM_A **238** to communicate with the VM_B **258**. In a third layer, a programmable packet switcher **236**, **238** in each network **230**, **250** exchanges packets between the networks **230**, **250** once the peering relationship is established.

(19) Each SDN controller **232**, **252** may be, for example, an application in the SDN architecture that manages flow control. For example, each SDN controller **232**, **252** may be a software system

or collection of systems that together provides management of network state, a high-level data model that captures the relationships between managed resources, policies and other services provided by the controller, an application programming interface (API) that exposes controller services to an application, etc. In some examples, the SDN controller may further provide a secure TCP control session between itself and associated agents in the network. Each SDN controller **232**, **252** may run on a server, and use protocols to tell the programmable packet switcher **236**, **238** where to send packets. The SDN controllers **232**, **252** may direct traffic according to one or more forwarding policies. An interface or protocol allows the SDN controllers **232**, **252** to communicate with other devices in the networks **230**, **250**. For example, the SDN controller **232** in the first network **230** may use such interface or protocol to communicate with the OHC **234**. The SDN controller **232** programs the OHC **234** to pull programming from the second SDN controller **252** in the second network **250**. Similarly, the second SDN controller **252** in the second network **250** programs the OHC **254** to pull programming from the first SDN controller **232** in the first network **230**. Such programming information pulled by the OHCs **234**, **254** is used to establish a peering relationship.

(20) Each OHC **234**, **254** may run on a same host as its associated VM **238**, **258**. For examples, first OHC **234** may run on a same first host as VM_A **238**, and second OHC **254** may run on a same second host as VM_B **258**. Each host may have an instance used to control the VMs **238**, **258**.

(21) Each OHC **234**, **254** may obtain information from the SDN controllers **232**, **252**. For example, the OHC **234**, **254** obtains first information from the SDN controller **232**, **252** in its own network. For example, such first information may include information about the VM **238**, **258** in the same network. As another example, such first information may include programming information from the SDN controller **232**, **252** in the same network, such programming information providing instructions for the OHC to pull second information from the SDN in another network. Such second information obtained from the SDN in the other network may include connectivity information to establish peering between the networks.

(22) According to some examples, the OHC **234**, **254** in one network may subscribe to the SDN controller **232**, **252** in the other network. For example, the first OHC **234** may subscribe with a local virtual machine controller to an endpoint meta-service, and receive updates about each VM_A **238** on the same host as the OHC **234**. For each endpoint received from the endpoint meta-service, the OHC **234** uses a corresponding identifier to subscribe to a peering service through which it receives updates on the second network **250**.

(23) When the first network **230** is attempting to establish a peering connection with the second network **250**, the first OHC **234** boosts only as much connectivity information from the second SDN controller **252** in the second network **250** as is needed to establish the connection. According to some examples, such connectivity information may include location information, load balancer information, transmittivity information, or any other information that may inform the VM_A **238** how to reach peer VM_B **258**. The location information may include, for example, an IP address, such as a destination address, forwarding information, node identifiers, etc. For example, the OHC **234** may use a corresponding endpoint identifier, peer network identifier, and host IP address to subscribe to destination and forwarding services for the second network **250**. When corresponding updates are received from the destination and forwarding services, the OHC **234** can use the updates to create a flow from VM_A **238** to VM_B **258**.

(24) The OHCs **234**, **254** push programming from the SDN controller **232**, **252** to the programmable packet switchers **236**, **256**. For example, when the first OHC **234** receives the first information from the SDN controller **232** in its network **230** and the second information from the SDN controller **252** in the second network **250**, it may push an aggregate of the first and second information to the packet switcher **236**. Upon receiving a packet from the VM_A **238**, the packet switcher **236** can directly forward the packet to the location of VM_B **258** without any proxy. By

way of example only, the programmable packet switcher can modify a header of the packet or encapsulate the packet with an additional header, wherein such header information for routing the packet directly to the VM_B was received from the SDN controller 252 via the OHC 234.

(25) The programmable packet switcher 236, 256 may be, for example, a data plane control, an application specific integrated circuit (ASIC), a network interface controller (NIC), a logical switch, or any other types of programmable switch. The programmable packet switcher 236, 256 may be adapted to execute any of a variety of programs, such as programs from the SDN controllers 232, 252 for sending and/or receiving packets between the networks 230, 250.

(26) Each of VM_A 238 and VM_B 258 may be a virtual environment that functions as a virtual computer system with its own central processing unit (CPU), memory, network interface, storage, etc. The first VM_A 238 of the first network 230 may share a host machine with the OHC 234. Similarly, the second VM_B 258 in the second network 250 may share a host machine with the OHC 254. In other examples, the VMs, OHCs, and/or SDNs may reside on any combination of independent host machines that are coupled to enable communication between the devices in each network.

(27) While two networks 230, 250 are shown, it should be understood that peering may be established among any number of networks. The peering may be a bidirectional peering relationship, wherein each network sees the other network and also shares its own resources with the other network. In other examples, the peering may be a one-way peering, such as where the first network peers into the second network, but does not allow the second network to peer into the first network. While a limited number of components are illustrated within each network 230, 250, it should be understood that additional components may be present. By way of example only, each network 230, 250 may include multiple VMs or other resources. As another example, each network may include a VPN, interconnect, or other components for establishing peering between the networks.

(28) The architecture described in connection with FIG. 2 provides a high bandwidth peering connection, as each VM can use a full egress bandwidth to communicate across the network. Accordingly, the networks can use full fabric cross-section bandwidth for communication. Moreover, because no proxies are needed to establish the communication between networks, the architecture is highly reliable.

(29) FIG. 3 illustrates an example of establishing a peering connection, including communications with OHCs 234, 258 in a first VPC and second VPC, respectively. OHC 234 may receive, from the first VPC, peering information identifying other VPCs for peering, such as VPC 2. For example, such peering information may include a virtual local area network extension identifier (VNID) for peer VPC2.

(30) The OHC 234 uses the received peering information to obtain connectivity information, such as destination and forwarding information, for the VM_B 258 in VPC2. For example, the OHC 234 may use the VNID for peer VPC2 to obtain such destination and forwarding information. According to some examples, the OHC 234 may subscribe to destination and/or forwarding services of the VPC2 for updates. The OHC 234 may receive such destination and forwarding information for each virtual machine or other endpoint in the VPC2. Using the obtained information, the OHC 234 can create a flow from VM_A 238 to VM_B 258. The flow may be created using, for example, a virtual internet protocol (IP) address of the VM_A 238, the VNID to which the VM_A 238 belongs, a physical IP address of a host that contains the VM_A 238, a virtual IP address of the VM_B 258, a peer VNID to which the VM_B 258 belongs, and a physical IP address of a host that contains VM2. A flow in the reverse direction, from VM_B 258 to VM_A 238, may be programmed similarly when the second OHC 254 obtains the peering information from the VPC2 and the connectivity information from the VPC1.

(31) FIG. 4 illustrates additional details of the OHC 234 in a first network, such as a first VPC, used for pairing with a second network, such as a second VPC.

(32) As shown, source VPC shard **420** includes peering provider **421**. Peering provider **421** may be, for example, a service or module that provides peering information to peering subscriber **441** of on-host controller (OHC) **440**. As described above, the peering information may include an identifier for other VPCs. In the present example, the peering provider **421** may provide an identifier for destination VPC shard **430** to the peering subscriber **441**.

(33) Using the peering information obtained by the peering subscriber **441**, the OHC **440** subscribes to the destination VPC shard **430**. As such, the OHC **440** may obtain connectivity information for each virtual machine or other endpoint in the destination VPC shard **430**, and use the connectivity information to establish a direct peering connection between a virtual machine or other endpoint in the source VPC shard **420** and each virtual machine or other endpoint in the destination VPC shard **430**. Peering destination provider **432** in the destination VPC shard **430** provides destination information to destination subscriber **442** in the OHC **440**. The information may be provided for each virtual machine in the destination VPC shard **430**. Peering destination provider **432** may be, for example, a routing table or module in communication with a routing table. The destination information provided to the destination subscriber **442** may be routing table information. Load balancer provider **433** in the destination VPC shard **430** provides load balancer information to load balancer subscriber **443** in the OHC **440**. Such load balancer information may include, for example, an indication of a current workload for a given virtual machine in the destination VPC, an indication of capacity of the given virtual machine, or other information. Forwarding provider **444** provides forwarding information to forwarding subscriber **444** in the OHC **440**.

(34) The subscribers **441**, **442**, **443**, **444** in the OHC **440** may each communicate with a corresponding set of providers **445**, **446**, **447**, **448**. According to some examples, the providers **445**, **446**, **447**, **448** may communicate with each other also. Peering provider **445** receives peering information from the peering subscriber **441**. Destination provider **446** in the OHC **440** receives the destination information from the destination subscriber **442** of the OHC **440** and provides it to a destination subscriber **462** in packet switcher **460**. Load balancer provider **447** in the OHC receives load balancing information from the load balancer subscriber **443** and provides it to load balancer subscriber **463** of the packet switcher **460**. Forwarding provider **448** receives the forwarding information from the forwarding subscriber **444** in the OHC **440** and provides it to forwarding subscriber **464** in the packet switcher **460**. The packet switcher **460** uses the information from the OHC providers **446**, **447**, **448** to establish a peering connection with individual virtual machines or other endpoints in the destination VPC shard **430**.

(35) FIG. 5 is a block diagram of an example environment for implementing the system described above. The system can be implemented on one or more devices having one or more processors in one or more locations, such as on one or more host machines **291**, **292**, **293** in a VPC **330**.

(36) The one or more host machines **291-293** in the VPC **330** may be coupled to one or more computing devices **512**, one or more storage devices **513**, and/or any of a variety of other types of devices over a network **515**. The storage device(s) **513** can be a combination of volatile and non-volatile memory, and can be at the same or different physical locations than the computing devices **512**. For example, the storage device(s) **513** can include any type of non-transitory computer readable medium capable of storing information, such as a hard-drive, solid state drive, tape drive, optical storage, memory card, ROM, RAM, DVD, CD-ROM, write-capable, and read-only memories.

(37) The storage device(s) **513**, in some examples, may be part of the VPC **330**. The VPC **330** can be configured to perform a number of cloud computing platform services using the host machines **291-293** and/or storage **513**, such as hosting cloud storage for data back-up, or hosting one or more virtual machines accessible by computing device **512** in communication therewith.

(38) Each host machine **291-293** can include one or more processors, memory, and/or other components typically found on host machines. Memory **392** can store information accessible by the

processor(s) **398**, including instructions **396** that can be executed by the processor(s) **398**. The memory **392** can also include data **394** that can be retrieved, manipulated or stored by the processor(s) **398**. The memory **392** can be a type of non-transitory computer readable medium capable of storing information accessible by the processor(s) **398**, such as volatile and non-volatile memory. The processor(s) **398** can include one or more central processing units (CPUs), graphic processing units (GPUs), field-programmable gate arrays (FPGAs), and/or application-specific integrated circuits (ASICs), such as tensor processing units (TPUs).

(39) The instructions **396** can include one or more instructions that when executed by the processor(s) **398**, cause the one or more processors to perform actions defined by the instructions. The instructions **396** can be stored in object code format for direct processing by the processor(s) **398**, or in other formats including interpretable scripts or collections of independent source code modules that are interpreted on demand or compiled in advance. The instructions **396** can include instructions for implementing the system consistent with aspects of this disclosure. Instructions **396** may further be executed to identify particular processes of virtual machines **238**.

(40) The data **394** can be retrieved, stored, or modified by the processor(s) **398** in accordance with the instructions **396**. The data **394** can be stored in computer registers, in a relational or non-relational database as a table having a plurality of different fields and records, or as JSON, YAML, proto, or XML documents. The data **394** can also be formatted in a computer-readable format such as, but not limited to, binary values, ASCII or Unicode. Moreover, the data **394** can include information sufficient to identify relevant information, such as numbers, descriptive text, proprietary codes, pointers, references to data stored in other memories, including other network locations, or information that is used by a function to calculate relevant data.

(41) Although FIG. 5 illustrates the processors **398** and the memory **392** as being within the host machine **292**, the processors **398** and the memories **392** can include multiple processors and memories that can operate in different physical housings and/or locations. For example, some of the instructions **396** and the data **394** can be stored on a removable SD card and others within a read-only computer chip. Some or all of the instructions and data can be stored in a location physically remote from, yet still accessible by, the processors **398**. Similarly, the processors **398** can include a collection of processors that can perform concurrent and/or sequential operation. The host machines **291-293** can each include one or more internal clocks providing timing information, which can be used for time measurement for operations and programs run by the host machines.

(42) Host machines **291, 293** may be configured similarly to the host machine **292**, including memory, processors, and other components. The computing device **512** can also include one or more processors, memory, and other components.

(43) The host machine **292** can be configured to receive requests to process data from the computing device **512**. For example, the VM_A **238** may provide a variety of services to users, through various user interfaces and/or APIs exposing the platform services. For example, one or more services can be a machine learning framework or a set of tools for managing software applications programmed to offer a particular service.

(44) The device **512** and the host machine **292** can be capable of direct and indirect communication over the network **515**. The devices can set up listening sockets that may accept an initiating connection for sending and receiving information. The network **515** itself can include various configurations and protocols including the Internet, World Wide Web, intranets, virtual private networks, wide area networks, local networks, and private networks using communication protocols proprietary to one or more companies. The network **515** can support a variety of short- and long-range connections. The short- and long-range connections may be made over different bandwidths, such as 2.402 GHz to 2.480 GHz 2.4 GHz and 5 GHz; or with a variety of communication standards, such as standards for wireless broadband communication. The network **515**, additionally or alternatively, can also support wired connections between the devices, including over various types of Ethernet connection.

(45) Although a single VPC and computing device 512 are shown in FIG. 5, aspects of the disclosure can be implemented according to a variety of different configurations and quantities of computing devices, including in paradigms for sequential or parallel processing, or over a distributed network of multiple devices. In some implementations, aspects of the disclosure can be performed on a single device, and any combination thereof.

(46) FIG. 6 is a block diagram illustrating connection between a peering group and a SDN data plane via a fully sharded SDN control plane. In this example, a resource model includes a plurality of networks, such as VPCs 1-*n*, in a peering group 600. In the peering group 600, each VPC 1-*n* is connected with each other. While several VPCs are shown, it should be understood that any number of VPCs may be included in the peering group 600.

(47) Because only a limited amount of information needed for peering between the VPCs is pulled, an SDN control plane may be sharded. In this example, a fully sharded SDN control plane includes a plurality of SDN controllers 232, 252, 272. Each SDN controller 232, 252, 272 is responsible for a subset of VPCs 1-*n*. For example, the peering group 600 may be sharded or partitioned such that each VPC 1-*n* is associated with one SDN controller 232, 252, 272. Each SDN controller 232, 252, 272 may control one, two, three, or any other number of VPCs 1-*n*. For example, as shown, a first SDN controller 232 controls VPC2 and another VPC, a second SDN controller 252 controls VPC1, VPC_{*n*}, and another VPC. SDN controller 272 controls VPC3, VPC4, and another VPC.

(48) Each SDN controller 232, 252, 272 may communicate with each of a plurality of OHCs 234, 244, 254, 264, 284, etc. in an SDN data plane. Each OHC may be associated with one of the VPCs in the peer group 600. For example, OHC 234 may be associated with VPC1, OHC 244 may be associated with VPC2, OHC 254 may be associated with VPC3, etc. The full mesh connectivity between each SDN controller 232, 252, 272 and each OHC 234-284 provides for the connectivity among each VPC 1-*n* in the peering group 600. For example, because each OHC 234-284 is coupled to each SDN controller 232, 252, 272, the OHC for any given VPC can pull connectivity information from a particular SDN controller that is responsible for another VPC with which the given VPC is to peer. By way of example, if VPC2 is to peer with VPC3, OHC 244 that is associated with VPC2 may pull connectivity information for VPC3 from SDN controller 272 to establish the peering relationship.

(49) The fully sharded SDN control plane provides scalability for connecting a large number of VPCs. By way of example only, thousands of peers or more may be connected. The SDN data plane, including the OHCs 234-284, pulls minimal programming information from the SDN controllers 232-272 on demand.

(50) FIG. 7 provides an example of route conflict resolution across peered VPCs. For example, all virtual machines using an on-premise router in a VPC may have IP addresses within a particular range, and a BGP announcement may be made to other devices to use such IP addresses to reach the virtual machines. Such routes may be received and managed using, for example, a route manager in an SDN controller. All routes learned from a VPC's on-premises router may be provided to a peer VPC, which can use the same addresses. In such instances, a conflict may arise between a route including the on-premises router for the first VPC and a route including an on-premises router for the peer VPC.

(51) If a route conflict arises, SDN controllers 732, 742, 752, 762 can communicate with each other to resolve the conflict. For example, each SDN controller 732, 742, 752, 762 may include a route manager. One or more or all of the SDN controllers 732, 742, 752, 762 may also include a peer route manager. For example, the peer route manager may be a separate module used for resolving conflicts, while the route managers are used to communicate the routes for reaching virtual machines or other endpoints in a VPC. One SDN controller may be designated for resolving the route conflict using the peer route manager. For example, the SDN controller that is responsible for programming VM egress may perform route conflict resolution. For example, the SDN controllers 732, 742, 752, 762 may exchange routes with each other over other services, such as by using the

route manager to communicate the route information to the peer route manager of the SDN controller responsible for the resolution, illustrated as SDN controller **732** in the present example. Such routing information communicated among the SDN controllers may include routing information from VPC**1** to complete a routing table, such as a destination address. The SDN controller **732** uses a pre-defined set of rules to determining a winning route when there is a conflict, and publishes the winning route to OHC **734**. In turn, the OHC **734** publishes the winning route, along with other relevant information pulled from the SDN controllers **732-762**, to programmable packet switcher **736**.

(52) According to some examples, the pre-defined set of rules for resolving route conflicts may relate to priority of routes. For example, one rule may be that routes in a local network have higher priority than routes in a peer network. This may ensure that local traffic is not redirected or taken away to a remote network. According to some examples, each peer network may have an associated preference, and routes may be assigned a priority based on that preference. In some examples, a value of an identifier may be used to determine priority, such as by assigning priority to a higher or lower number identifier. According to further examples, priority may be assigned based on a routing policy associated with a peer network.

(53) FIG. **8** illustrates an example of scalable cross-VPC connectivity via gateways. For example, one or more scalable gateways **840** may be implemented between programmable packet switcher **236** of a first VPC and programmable packet switcher **256** of a second VPC. When the programmable packet switch **236** receives a packet from VM_A **238**, the packet switcher **236** may load balance the packet to the one or more scalable gateways **840**. The one or more scalable gateways **840** may have full network state, and forward the packets to a destination, such as programmable packet switcher **256**. For example, the scalable gateways **840** may see every virtual machine, load balancer, or other endpoint in each VPC.

(54) According to some examples, the one or more scalable gateways **840** may be a pool of scalable gateways. Each scalable gateway may be, for example, a dedicated hardware unit, a virtual machine, virtual router, or other gateway.

(55) According to some examples, the scalable gateways **840** may be used to shift traffic from a first VPC to one or more second VPCs when the first VPC becomes overloaded. For example, the scalable gateways **840** may monitor traffic demands and capacity within the first VPC, as well as for the one or more second VPCs. When a load on VM_A **238** in the first VPC becomes too high relative to its capacity, the scalable gateway may shift traffic to one or more second VPCs using a peering connection established by the OHC **234** of the first VPC. According to some examples, the scalable gateways **840** may prompt the OHC **234** to initiate the peering and/or to request connectivity information from the SDN controller **252** of the second VPC. Accordingly, the VM_B **258** of the second VPC may share some of the load of the VPC_A **238**.

(56) FIG. **9** illustrates an example of cross-network multiplexing. Each VPC publishes resources **950**, such as resources from load balancers, virtual machines, BGP routes, peering, etc., to a peering multiplexer **910** in SDN controller **902**. The peering multiplexer **910** uses such published resources to derive a virtual network state **930**. The virtual network state **930** may indicate all resources in a network, such as all virtual machines across all VPCs, all load balancers across all VPCs, all BGP routes across all VPCs, etc.

(57) The peering multiplexer **910** in SDN controller **902** may be used to apply peering policies when processing updates for endpoints, ILBs, subnets and routes. The peering multiplexer **610** may, for each resource, keep track of which peering-related services need to be updated. Moreover, for each peering policy, the peering multiplexer **910** may create peering-related services by pulling resource state from resource-managers. Depending on the policy, the peering multiplexer **910** may filter out programming on peering-related services.

(58) Peering manager **920** maintains a map of all peering policies that are relevant to the VPC. For example, for each peering-policy, it maintains a list of identifiers of peer networks that are using

the policy. The policies may include, for example, filtering policies or other types of policies. The filtering policies may determine whether particular protocols are allowed, whether particular routes are allowed, whether particular addresses or subnets are allowed, or any combination of these or other parameters, etc. Each pair of VPC pairings may have an independent policy, referred to as active export policies, which may be managed by peer manager **920**.

(59) The SDN controller **902** can publish information on routing to resources such as VMs and load balancers. The SDN controller **902** can also publish information on border gateway protocol (BGP) routes for federation. Federation refers to all SDNs communicating with each other to exchange routes and resolve route conflicts. For example, the peer multiplexer **910** may provide federated route exchange information, such as when communicating between a route manager and a peer route manager as illustrated in FIG. 7. Peer multiplexer **910** may provide policies to routing service, such as the SDN controller providing peering information to the OHC.

(60) Different services may be published per connectivity policy per VPC. For example, different services may be published for an IPv6 only policy, a specific exported LP ranges policy, a private addresses only policy, etc. A resource can be published on multiple services simultaneously. Multiplexing may be performed on the resources based on export policies for a VPC.

(61) The SDN controller **901** may provide for optimized policy based filtering. The SDN controller **901** may determine the set of unique active peering exchange policies for each VPC using peering manager **920**. Peering multiplexer **910** then only publishes the set of resources for each unique exchange policy per VPC. The SDN controller **902** and programmable hosts responsible for a peered VPC then subscribe to the appropriate set of resources based on an exchange policy identifier. This multiplexer architecture results in a much smaller set of services published as compared to an architecture in which a different set of resources is published between each pair of peered VPC.

(62) FIG. 10 illustrates an example method **1000** of establishing a peering connection between VPCs. While operations are described in a particular order, the order may be modified. In some examples, operations may be performed in parallel. In some examples, operations may be added or omitted. The operations are described in the method **1000** from a perspective of an OHC in a source network in establishing peering with a target network. Corresponding operations may be performed by other components, such as SDN controllers, packet switchers, and the like.

(63) In block **1010**, the OHC in a source VPC receives programming information from a SDN controller within the source VPC. The programming information may include, for example, an identification of peer networks, policy identifiers, and information for use by the source OHC in reaching a target SDN controller of the target VPC.

(64) In block **1020**, the source OHC receives connectivity information from the target SDN controller in the target VPC. For example, such connectivity information may include a destination address for the VPC, such as an IP address for a target virtual machine in the target VPC. Such connectivity information may include forwarding information for sending packets to the target VPC. The source OHC may subscribe to a service of the target VPC and thereby receive updates in connectivity information from the target VPC.

(65) In block **1030**, the source OHC controller establishes a peering connection from the source VPC to the target VPC using the received connectivity information. Such peering connection may be established based on the limited connectivity information received, without requiring a copy of the entire target VPC. Moreover, such peering connection may be established without any proxies.

(66) In block **1040**, the source OHC provides packet information to the packet switcher in the source VPC. Such packet information may be used for direct communication between a first virtual machine in the source VPC and a second virtual machine in the target VPC. The packet information may include, for example, destination, forwarding, and authentication information received from the target SDN controller. Accordingly, the source packet switcher may exchange packets directly with the target VPC, without proxies.

(67) Unless otherwise stated, the foregoing alternative examples are not mutually exclusive, but may be implemented in various combinations to achieve unique advantages. As these and other variations and combinations of the features discussed above can be utilized without departing from the subject matter defined by the claims, the foregoing description of the embodiments should be taken by way of illustration rather than by way of limitation of the subject matter defined by the claims. In addition, the provision of the examples described herein, as well as clauses phrased as “such as,” “including” and the like, should not be interpreted as limiting the subject matter of the claims to the specific examples; rather, the examples are intended to illustrate only one of many possible embodiments. Further, the same reference numbers in different drawings can identify the same or similar elements.

Claims

1. A system, comprising: a software-defined networking (SDN) controller; an on-host controller; and a programmable packet switcher; wherein the SDN controller, on-host controller, and programmable packet switcher are communicatively coupled within a first network; wherein the on-host controller in the first network receives programming information from the SDN controller in the first network, the programming information providing instructions for the on-host controller to pull limited connectivity information directly from a second SDN controller in a target network, and use the limited connectivity information to establish a proxyless peering connection between the first network and the target network.
2. The system of claim 1, further comprising a first virtual machine within the first network.
3. The system of claim 2, wherein the first virtual machine and the on-host controller reside on a same host machine.
4. The system of claim 1, wherein the first network is a virtual private cloud (VPC).
5. The system of claim 1, wherein the limited connectivity information comprises a destination address in the target network.
6. The system of claim 5, wherein the limited connectivity information further comprises forwarding information.
7. The system of claim 1, wherein the programmable packet switcher delivers packets from a first virtual machine in the first network to a second virtual machine in the target network using the limited connectivity information.
8. The system of claim 1, wherein the first network includes a plurality of endpoints, and wherein each endpoint in the source network is adapted to communicate directly with each endpoint in the target network using the proxyless peering connection.
9. The system of claim 1, further comprising a plurality of SDN controllers in a fully sharded control plane, each of the plurality of SDN controller responsible for a subset of virtual private cloud (VPC) networks in a peering group, each VPC network having an associated on-host controller, wherein each of the plurality of SDN controllers is coupled to each of the associated on-host controllers.
10. The system of claim 1, further comprising a scalable gateway between the programmable packet switcher of the first network and a second programmable packet switcher of the target network.
11. A method of establishing a peering connection between a first network and a second network, the method comprising: receiving, at an on-host controller in the first network, programming information from a first controller within the first network, the programming information providing instructions for the on-host controller to pull connectivity information directly from a second controller within the second network; pulling, at the on-host controller in the first network, connectivity information directly from the second controller within the second network; establishing a proxyless peering connection between the first network and the second network

using at least the connectivity information; and providing, to a packet switcher within the first network, packet information for direct communication between a first endpoint in the first network and a second endpoint in the second network.

12. The method of claim 11, wherein the first endpoint is a virtual machine.

13. The method of claim 11, wherein receiving the connectivity information from the second controller comprises using the programming information from the first controller to subscribe to the second controller.

14. The method of claim 11, wherein the connectivity information comprises a destination address in the second network and forwarding information.

15. The method of claim 11, further comprising delivering, by a programmable packet switcher, packets from a first virtual machine in the first network to a second virtual machine in the second network using the connectivity information.

16. A software-defined networking (SDN) controller configured to establish a peering connection between a first network that includes the SDN controller and a second network, the SDN controller comprising: memory; and one or more processors in communication with the memory, the one or more processors configured to: provide programming information to an on-host controller in the first network, the programming information identifying the second network, the programming information providing instructions for the on-host controller to pull limited connectivity information directly from a second SDN controller; and cause the on-host controller to subscribe to the second SDN controller in the second network and pull the limited connectivity information directly from the second SDN controller.

17. The SDN controller of claim 16, comprising a peering multiplexer configured to receive information from a plurality of resources and publish routing information for communicating with the resources.

18. The SDN controller of claim 17, comprising a peering manager maintaining one or more export policies for the first network, wherein the peering multiplexer performs multiplexing on the resources based on the export policies.

19. The SDN controller of claim 16, comprising a routing manager configured to receive route information from one or more other SDN controllers.

20. The SDN controller of claim 19, wherein the SDN controller is configured to identify route conflicts based on the route information received and to resolve the conflict using a pre-defined set of rules.
