

(12) **United States Patent**
Cook

(10) **Patent No.:** **US 12,388,725 B2**
(45) **Date of Patent:** ***Aug. 12, 2025**

(54) **RESOURCE MEASUREMENT AND
MANAGEMENT**

(71) Applicant: **Comcast Cable Communications,
LLC, Philadelphia, PA (US)**

(72) Inventor: **Michael J. Cook, Flemington, NJ (US)**

(73) Assignee: **Comcast Cable Communications,
LLC, Philadelphia, PA (US)**

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 0 days.

This patent is subject to a terminal dis-
claimer.

(21) Appl. No.: **18/654,177**

(22) Filed: **May 3, 2024**

(65) **Prior Publication Data**

US 2024/0297831 A1 Sep. 5, 2024

Related U.S. Application Data

(63) Continuation of application No. 18/346,393, filed on
Jul. 3, 2023, now Pat. No. 12,003,389, which is a
continuation of application No. 17/712,858, filed on
Apr. 4, 2022, now Pat. No. 11,736,369, which is a
continuation of application No. 13/245,897, filed on
Sep. 27, 2011, now Pat. No. 11,323,337.

(51) **Int. Cl.**

H04L 41/5025 (2022.01)
G06Q 30/02 (2023.01)
H04L 43/0882 (2022.01)
H04L 47/26 (2022.01)
H04N 21/24 (2011.01)
H04N 21/2543 (2011.01)
H04N 21/647 (2011.01)

(52) **U.S. Cl.**

CPC **H04L 41/5025** (2013.01); **G06Q 30/02**
(2013.01); **H04L 43/0882** (2013.01); **H04L**
47/26 (2013.01); **H04N 21/2408** (2013.01);
H04N 21/2543 (2013.01); **H04N 21/64723**
(2013.01)

(58) **Field of Classification Search**

CPC . H04L 41/5025; H04L 43/0882; H04L 47/26;
G06Q 30/02; H04N 21/2408; H04N
21/2543; H04N 21/64723

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,361,394 A 11/1994 Shigihara
5,687,167 A 11/1997 Bertin et al.
5,761,191 A 6/1998 VanDervort et al.
5,872,588 A 2/1999 Aras et al.
5,920,701 A 7/1999 Miller et al.

(Continued)

Primary Examiner — Nazia Naoreen


(74) *Attorney, Agent, or Firm* — Banner & Witcoff, Ltd.

(57)

ABSTRACT

Aspects of the disclosure relate to measuring and managing
data traffic in one or more networks. In some embodiments,
a monitor may measure the traffic at one or more locations
within the network(s) or devices associated therewith to
determine whether the traffic exceeds a threshold. When the
traffic exceeds the threshold, one or more actions may be
taken, such as issuing or transmitting a command or direc-
tive. The command or directive may advise a device or an
application to throttle or reduce an input or stimulus respon-
sible for generating the traffic. In some embodiments, a
throttling may be effectuated to reduce the data traffic.

24 Claims, 6 Drawing Sheets



(56)

References Cited

U.S. PATENT DOCUMENTS

5,926,476	A	7/1999	Ghaibeh	2005/0195968	A1	9/2005	Park et al.
6,028,860	A	2/2000	Laubach et al.	2005/0281235	A1	12/2005	Beach et al.
6,125,105	A	9/2000	Edwards et al.	2006/0053146	A1	3/2006	Allhusen et al.
6,148,081	A	11/2000	Szymanski et al.	2006/0077960	A1	4/2006	Chang et al.
6,331,987	B1	12/2001	Beser	2006/0099956	A1	5/2006	Harada et al.
6,427,238	B1	7/2002	Goodman et al.	2006/0120282	A1	6/2006	Carlson et al.
6,502,135	B1	12/2002	Munger et al.	2006/0129458	A1*	6/2006	Maggio G06Q 30/0218
6,560,243	B1	5/2003	Mogul				705/14.2
6,594,246	B1	7/2003	Jorgensen	2006/0142040	A1	6/2006	Jones
6,915,530	B1	7/2005	Kauffman et al.	2006/0184664	A1	8/2006	Jung
6,941,574	B1	9/2005	Broadwin et al.	2006/0212917	A1	9/2006	Boucher et al.
6,970,641	B1	11/2005	Pierre et al.	2006/0233108	A1	10/2006	Krishnan
7,017,175	B2	3/2006	Alao et al.	2007/0071026	A1	3/2007	Rogers
7,027,461	B1	4/2006	Bontempi	2007/0078707	A1*	4/2007	Axe G06Q 30/02
7,055,169	B2	5/2006	Delpuch et al.				705/14.69
7,065,780	B2	6/2006	Barbier et al.	2007/0091176	A1	4/2007	Shih et al.
7,231,455	B2	6/2007	Marejka et al.	2007/0195789	A1	8/2007	Yao
7,305,697	B2	12/2007	Alao et al.	2007/0214248	A1	9/2007	Ikeda et al.
7,362,704	B2	4/2008	Sisto et al.	2008/0075003	A1	3/2008	Lee et al.
7,426,185	B1	9/2008	Musacchio et al.	2008/0089240	A1	4/2008	Toomey et al.
7,433,953	B1	10/2008	Kappler et al.	2008/0144586	A1	6/2008	Knecht et al.
7,440,999	B2*	10/2008	DeCenzo H04N 21/6587	2008/0191816	A1	8/2008	Balachandran et al.
			725/101	2008/0310452	A1	12/2008	Vedantham et al.
7,468,945	B2	12/2008	Enomoto et al.	2009/0067328	A1	3/2009	Morris et al.
7,650,379	B2	1/2010	Hart et al.	2009/0086634	A1	4/2009	Kwan et al.
7,876,696	B2*	1/2011	Kotelba H04L 1/0002	2009/0089825	A1*	4/2009	Coldwell H04N 7/17318
			370/468				725/25
7,917,648	B2	3/2011	Wu et al.	2009/0175235	A1	7/2009	Spinar et al.
8,201,205	B2	6/2012	Boucher et al.	2009/0213871	A1	8/2009	Carlson et al.
8,208,001	B2*	6/2012	Tucker H04N 21/6125	2009/0222851	A1*	9/2009	Talmi H04N 7/17336
			348/14.09				725/32
8,213,301	B2	7/2012	Park et al.	2009/0287405	A1	11/2009	Liu et al.
8,396,055	B2	3/2013	Patel et al.	2010/0014423	A1	1/2010	Furuskar et al.
8,682,363	B1	3/2014	Cardi et al.	2010/0085923	A1	4/2010	Nyberg et al.
8,843,652	B2	9/2014	Balasubramanian	2011/0016213	A1	1/2011	Nakao
8,850,480	B2	9/2014	Chane et al.	2011/0122887	A1	5/2011	Ma et al.
8,914,839	B2*	12/2014	Howe H04N 21/8166	2011/0134746	A1	6/2011	Liu et al.
			725/100	2011/0161261	A1	6/2011	Wu et al.
9,106,557	B2	8/2015	Wolcott et al.	2011/0194538	A1	8/2011	Zheng et al.
9,461,955	B1	10/2016	Sherrets	2011/0305451	A1	12/2011	Nishitani et al.
9,635,421	B2	4/2017	Tidwell et al.	2011/0306351	A1	12/2011	Pettersson et al.
10,735,347	B2	8/2020	Boucher et al.	2012/0039169	A1	2/2012	Susitaival et al.
2002/0006111	A1	1/2002	Akita et al.	2012/0057478	A1	3/2012	Brannstrom et al.
2002/0126673	A1	9/2002	Dagli et al.	2012/0063377	A1	3/2012	Osterling et al.
2002/0142780	A1	10/2002	Airy et al.	2012/0064908	A1	3/2012	Fox et al.
2002/0188732	A1	12/2002	Buckman et al.	2012/0099430	A1	4/2012	Vos et al.
2003/0018745	A1*	1/2003	McGowan H04L 65/611	2012/0116863	A1	5/2012	Boss et al.
			348/E7.071	2012/0203822	A1	8/2012	Floyd et al.
2003/0210665	A1	11/2003	Salmenkaita et al.	2012/0330710	A1	12/2012	Hauser et al.
2004/0177087	A1	9/2004	Wu et al.	2013/0016675	A1	1/2013	Cha et al.
2004/0257995	A1	12/2004	Sandy et al.	2013/0051359	A1	2/2013	Skarby
2005/0047783	A1	3/2005	Sisto et al.	2013/0121147	A1	5/2013	Tapia et al.
2005/0076339	A1	4/2005	Merril et al.	2013/0166773	A1	6/2013	Armstrong et al.
2005/0094558	A1	5/2005	Lu	2013/0201996	A1	8/2013	Masputra et al.
2005/0138674	A1*	6/2005	Howard H04N 21/8547	2014/0146666	A1	5/2014	Kwan et al.
			725/135	2014/0173025	A1	6/2014	Killick
				2014/0180575	A1	6/2014	Lin
				2014/0269962	A1	9/2014	Gao et al.

* cited by examiner

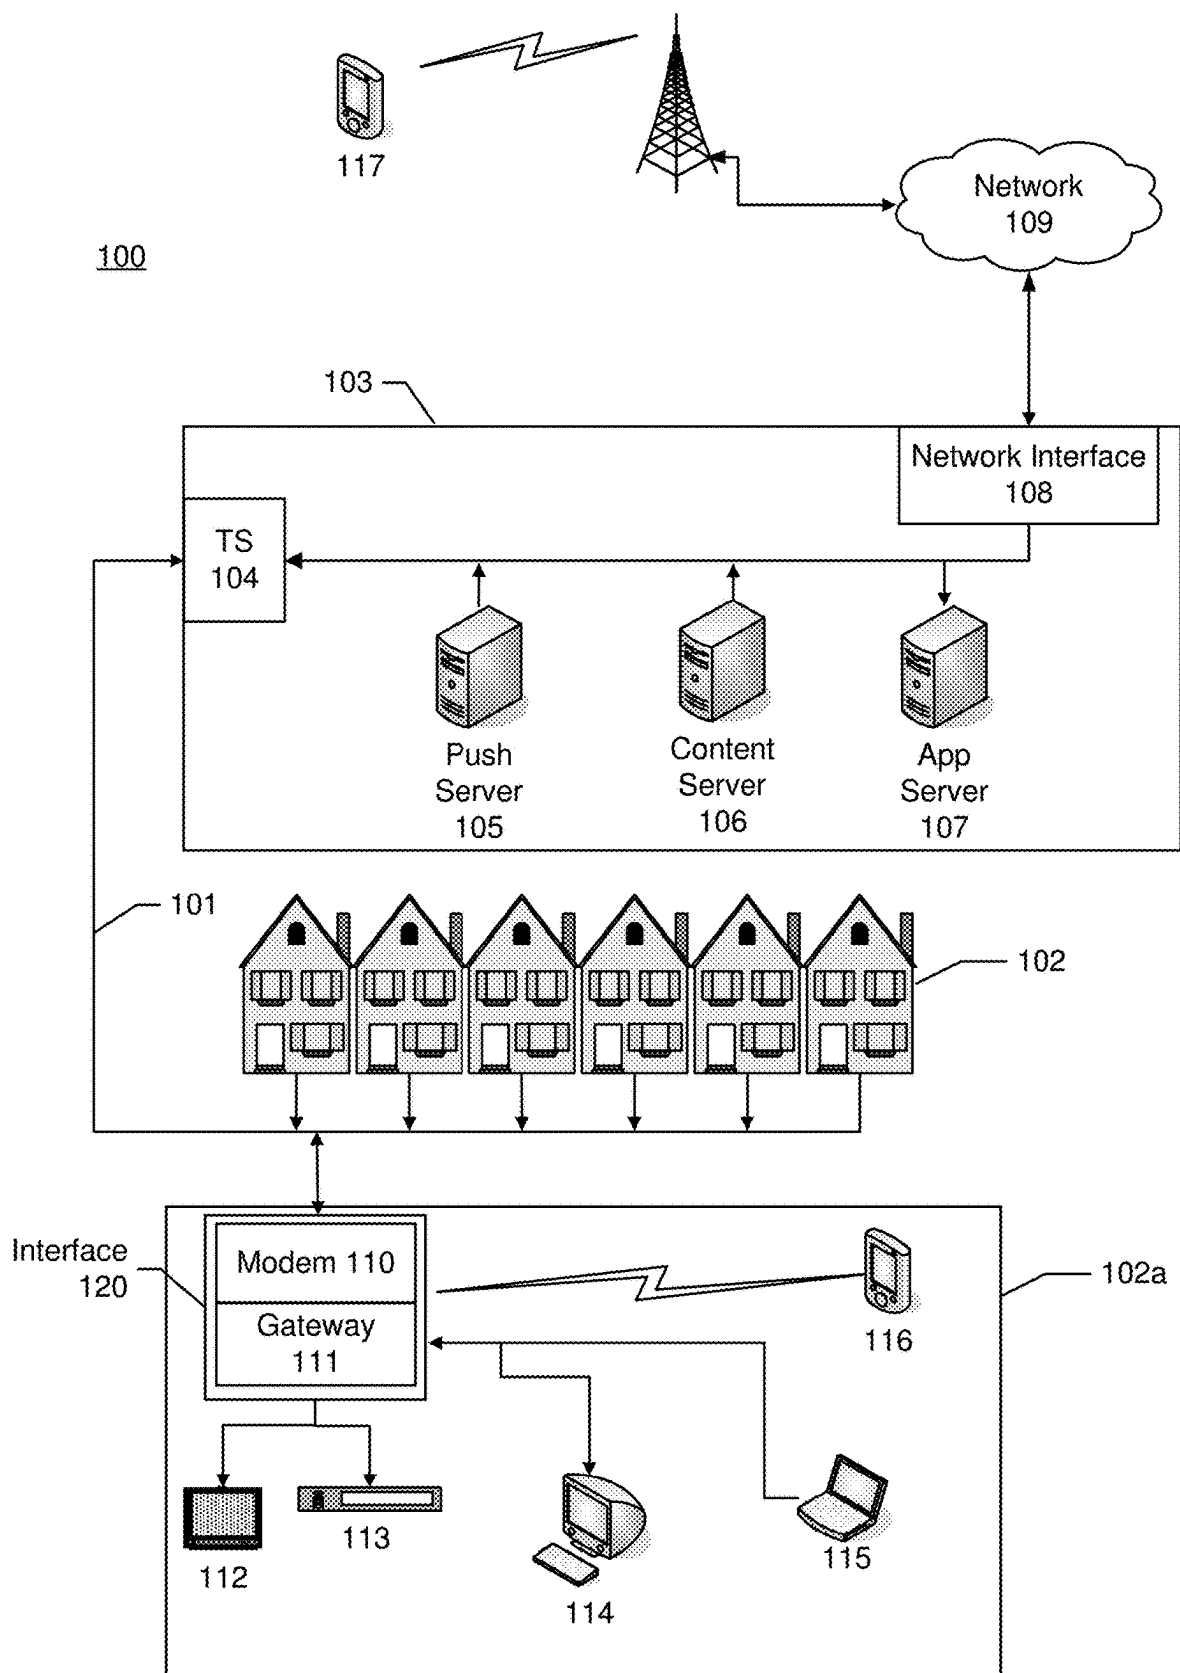


FIG. 1

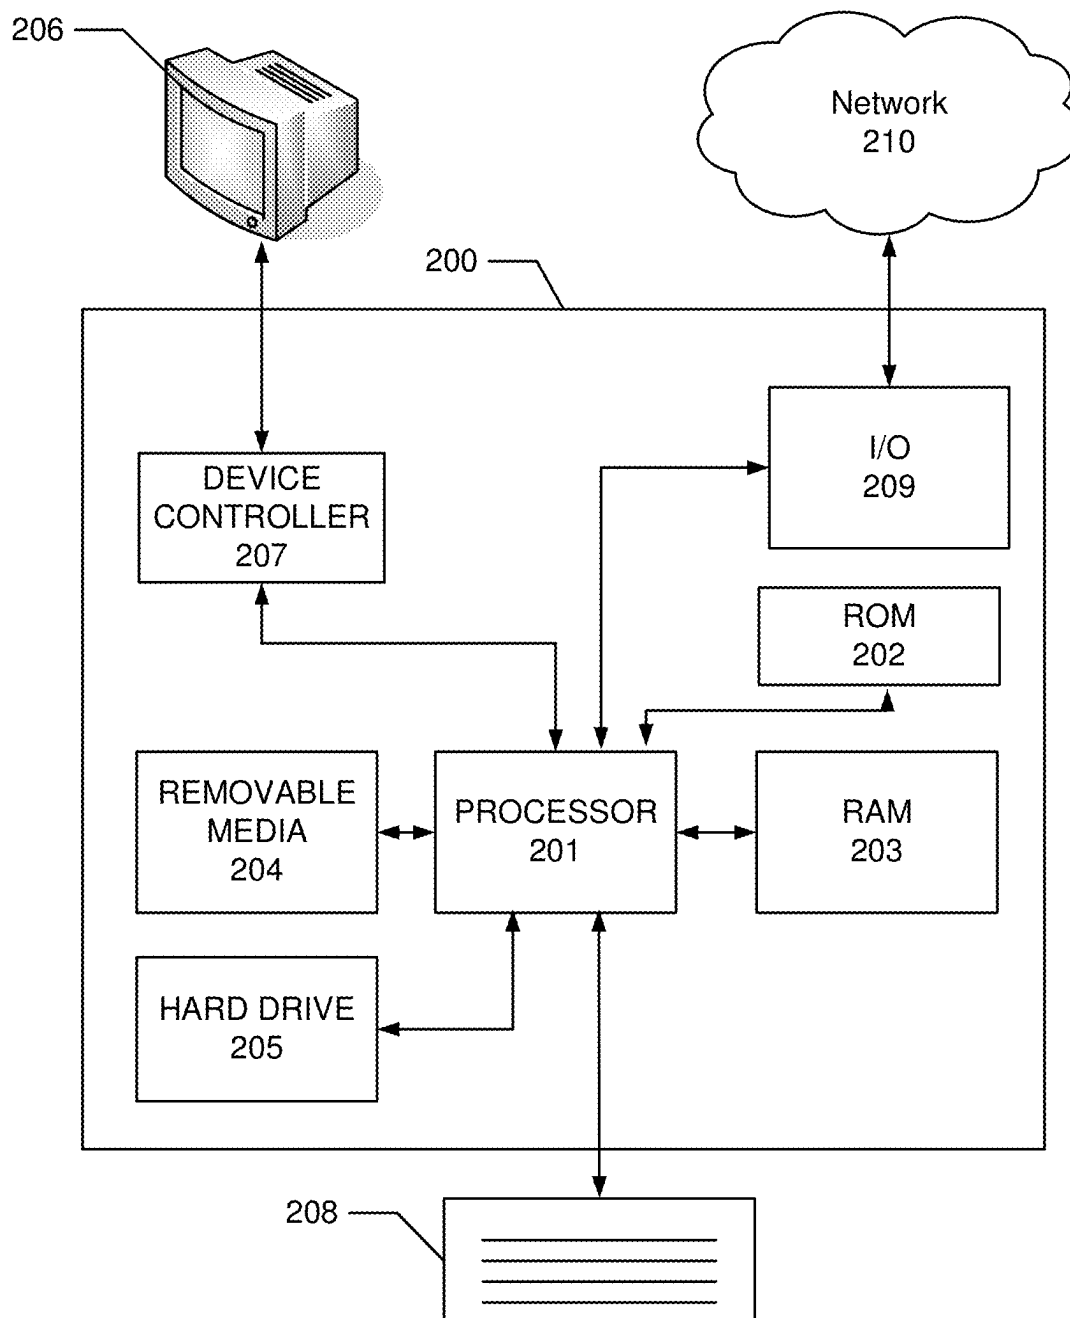


FIG. 2

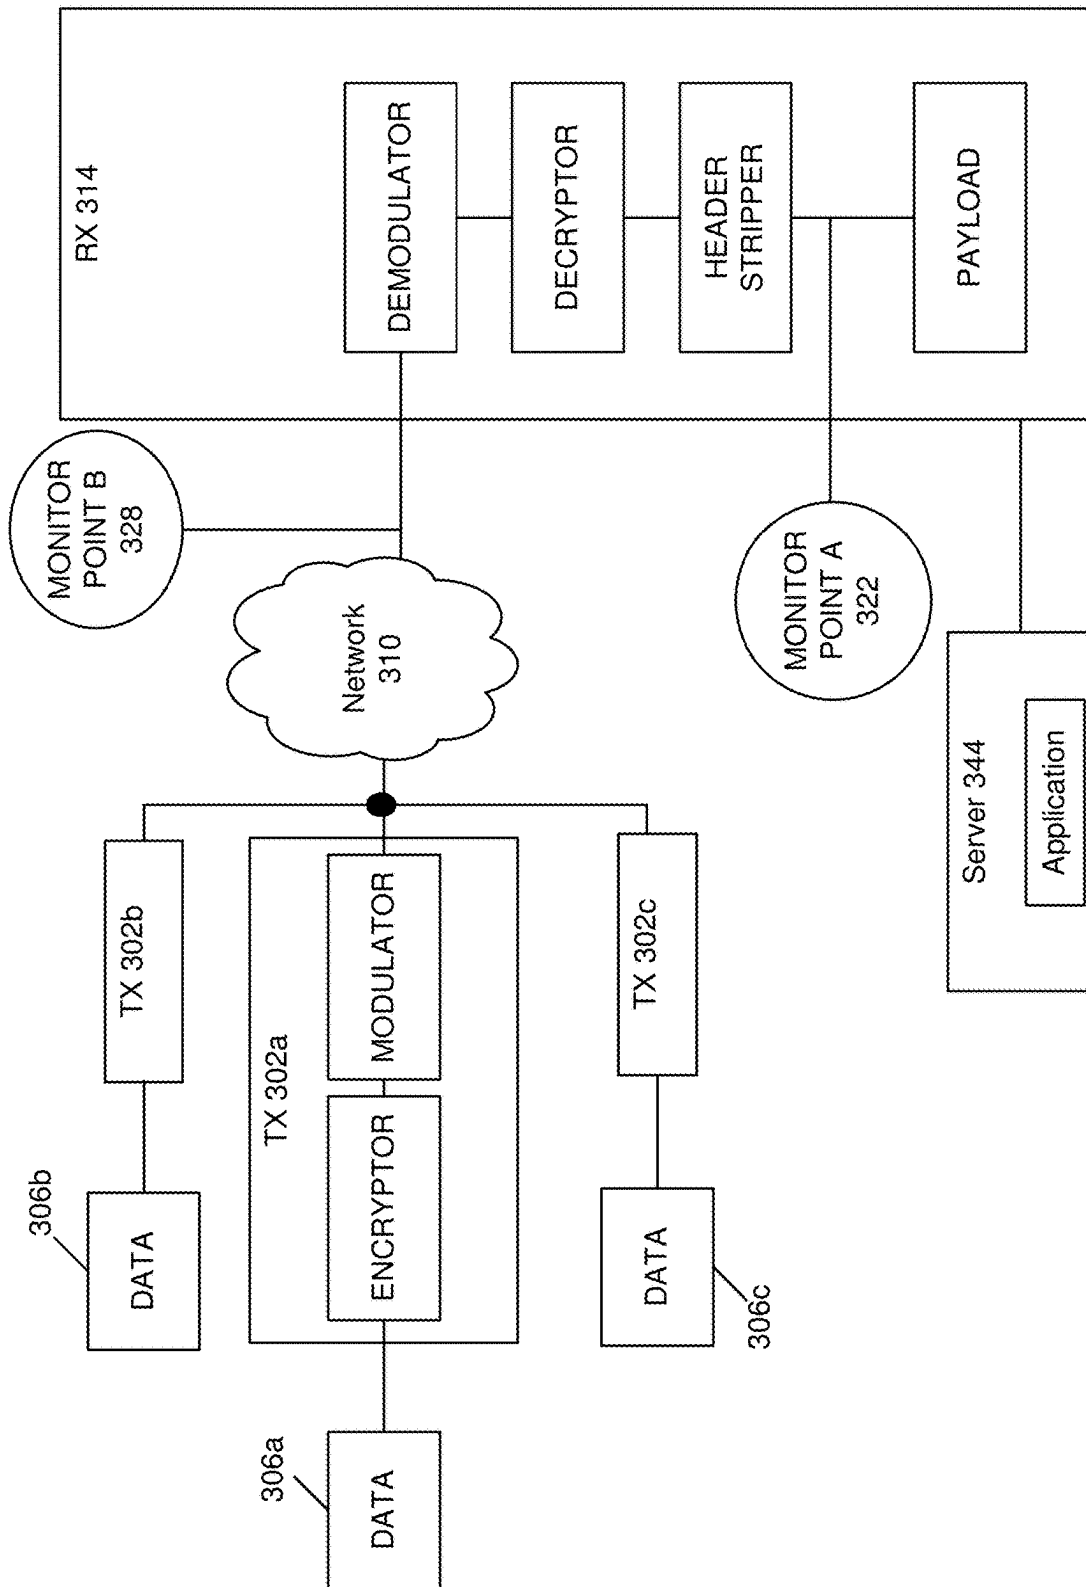


FIG. 3




FIG. 4




FIG. 5

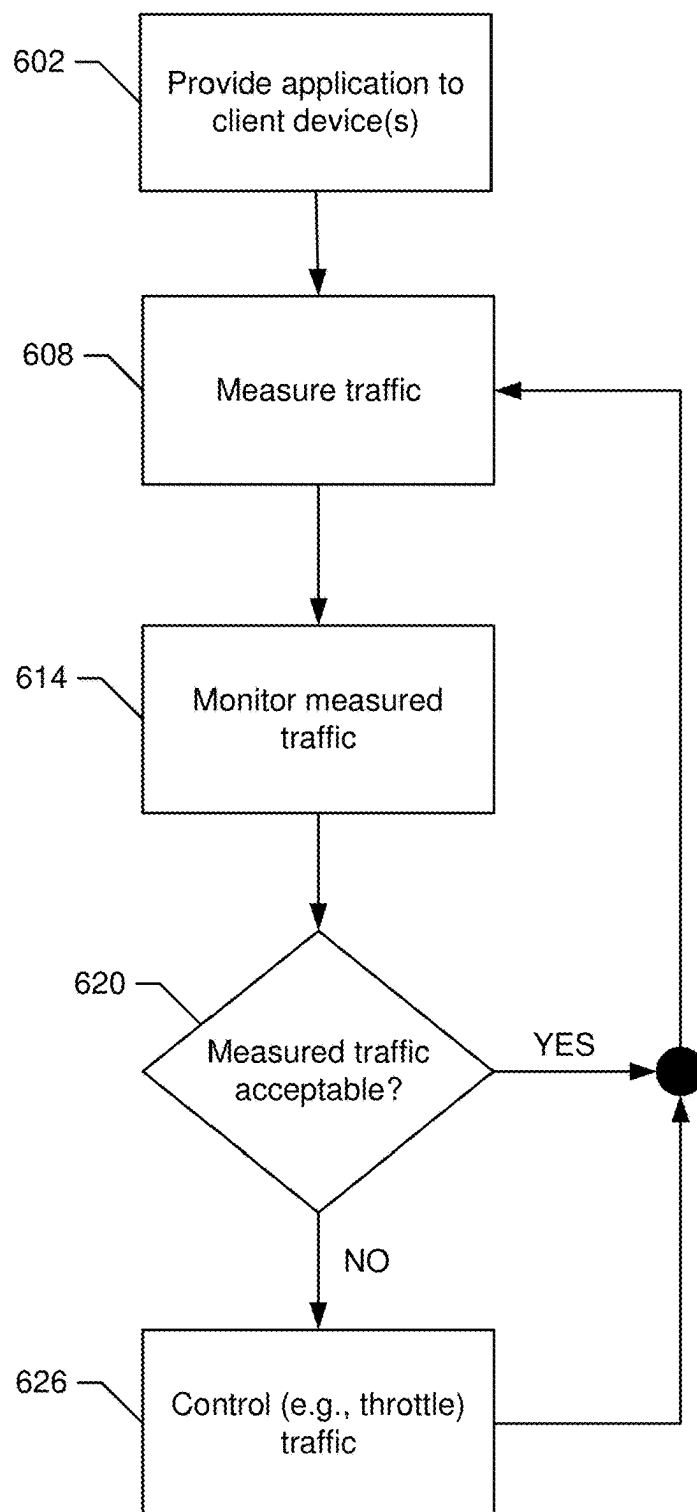


FIG. 6

1

RESOURCE MEASUREMENT AND MANAGEMENT

CROSS-REFERENCE TO RELATED APPLICATIONS

This application is a continuation of U.S. patent application Ser. No. 18/346,393, filed Jul. 3, 2023, which is a continuation of U.S. patent application Ser. No. 17/712,858 (now U.S. Pat. No. 11,736,369), filed Apr. 4, 2022, which is a continuation of U.S. patent application Ser. No. 13/245,897, filed Sep. 27, 2011 (now U.S. Pat. No. 11,323,337). Each of the above-referenced applications is hereby incorporated by reference in its entirety.

BACKGROUND

The transmission and distribution of data and content over networks has evolved over the years. Users expect a network, service, or content provider to be able to deliver data and content in various formats and at various times.

For particular system configurations and associated resources (e.g., bandwidth), dedicating resources to particular users may be expensive and burdensome. For example, in the context of an interactive environment (e.g., an interactive program guide (IPG)), offering users an ability to interact with content may deplete resources (e.g., bandwidth). The reduction or depletion of resources may have a noticeable impact or effect. For example, services might not be available as a result of the reduction in resources.

In the context of contention based protocols, frames of data may be sent by a transmitter when a frame becomes available. A frame may be resent until a receiver successfully receives it. As a network becomes more complicated (e.g., Ethernet systems with multiple sources and destinations) and/or communication volume increases, data frames may collide at a greater rate. An increase in collisions may result in a reduction or degradation in terms of efficiency.

Increasing the total amount of bandwidth available is expensive, as maintenance costs, management costs, and network complexity tend to increase with the increase in bandwidth, and resources may tend to be underutilized (on average) in such a configuration. Reducing the amount of bandwidth consumed by a given application may be undesirable, as an author or an operator of a given network or service may strive to provide content-rich applications.

SUMMARY

This summary is not intended to identify critical or essential features of the disclosure provided herein, but instead merely summarizes certain features and variations thereof.

In some illustrative embodiments, a monitor may be used to observe traffic or data flow in one or more networks. In some embodiments, the monitor may comprise one or more sniffers. In some embodiments, the monitor may observe the traffic in real-time or substantially real-time. In some embodiments, the monitor may observe traffic in an upstream direction, a downstream direction, or in both upstream and downstream directions.

In some embodiments, the monitor may be positioned at one or more locations. For example, the monitor may be positioned at an interface between one or more networks (e.g., at an interface between a radio frequency (RF) network and an internet protocol (IP) network). In some embodiments, the monitor may be positioned within one or more

2

networks (e.g., within an RF network). In some embodiments, the monitor may be positioned in, or located within, one or more devices. In some embodiments, the monitor may be positioned to receive results (e.g., traffic results) from one or more devices (e.g., one or more servers, such as one or more application servers). The monitor may analyze the received results and provide commands or directives to the one or more devices based at least in part on the analysis.

In some embodiments, the monitor may be used to manage or shape traffic in one or more networks. For example, in some embodiments the monitor may transmit or issue commands or directives to control the rate of traffic or data flow in the one or more networks.

Other details and features will also be described in the sections that follow.

BRIEF DESCRIPTION OF THE DRAWINGS

The present disclosure is pointed out with particularity in the appended claims. Features of the disclosure will become more apparent upon a review of this disclosure in its entirety, including the drawing figures provided herewith.

Some features herein are illustrated by way of example, and not by way of limitation, in the figures of the accompanying drawings and in which like reference numerals refer to similar elements.

FIG. 1 illustrates an exemplary information distribution network.

FIG. 2 illustrates an exemplary hardware platform on which various elements described herein can be implemented.

FIGS. 3-4 illustrate environments for measuring and monitoring traffic in which one or more aspects of this disclosure may be implemented.

FIG. 5 illustrates an exemplary monitor in accordance with one or more aspects of this disclosure.

FIG. 6 illustrates a method for demonstrating one or more aspects of this disclosure.

DETAILED DESCRIPTION

Various connections and/or communication links between elements are discussed in the following description. These connections and/or links are general and, unless specified otherwise, may be direct or indirect, wired or wireless, and this specification is not intended to be limiting in this respect.

In the following description of various illustrative embodiments, reference is made to the accompanying drawings, which form a part hereof, and in which is shown, by way of illustration, various embodiments in which aspects of the disclosure may be practiced. It is to be understood that other embodiments may be utilized and structural and functional modifications may be made, without departing from the scope of the present disclosure.

FIG. 1 illustrates an example information distribution network **100** on which many of the various features described herein may be implemented. Network **100** may be any type of information distribution network, such as satellite, telephone, cellular, wireless, etc. One example may be an optical fiber network, a coaxial cable network or a hybrid fiber/coax distribution network. Such networks **100** use a series of interconnected communication lines **101** (e.g., coaxial cables, optical fibers, wireless, etc.) to connect multiple premises **102** (e.g., businesses, homes, consumer dwellings, etc.) to a central office or headend **103**. The central office **103** may transmit downstream information

3

signals onto the lines **101**, and each premises **102** may have a receiver used to receive and process those signals.

There may be one line **101** originating from the central office **103**, and it may be split a number of times to distribute the signal to various premises **102** in the vicinity (which may be many miles) of the central office **103**. The lines **101** may include components not illustrated, such as splitters, filters, amplifiers, etc. to help convey the signal clearly, but in general each split introduces a bit of signal degradation. Portions of the lines **101** may also be implemented with fiber-optic cable, while other portions may be implemented with coaxial cable, other lines, or wireless communication paths. By running fiber optic cable along some portions, for example, signal degradation in those portions may be significantly minimized, allowing a single central office **103** to reach even farther with its network of lines **101** than before.

The central office **103** may include a modem termination system (TS) **104**, such as a cable modem termination system (CMTS), which may be a computing device configured to manage communications between devices on the network of lines **101** and backend devices such as servers **105-107** (to be discussed further below). The termination system (TS) may be as specified in a standard, such as the Data Over Cable Service Interface Specification (DOCSIS) standard, published by Cable Television Laboratories, Inc. (a.k.a. CableLabs), or it may be a similar or modified device instead. The termination system (TS) may be configured to place data on one or more downstream frequencies to be received by modems at the various premises **102**, and to receive upstream communications from those modems on one or more upstream frequencies. The central office **103** may also include one or more network interfaces **108**, which can permit the central office **103** to communicate with various other external networks **109**. These networks **109** may include, for example, networks of Internet devices, telephone networks, cellular telephone networks, fiber optic networks, local wireless networks (e.g., WiMAX), satellite networks, and any other desired network, and the interface **108** may include the corresponding circuitry needed to communicate on the network **109**, and to other devices on the network such as a cellular telephone network and its corresponding cell phones.

As noted above, the central office **103** may include a variety of servers **105-107** that may be configured to perform various functions. For example, the central office **103** may include a push notification server **105**. The push notification server **105** may generate push notifications to deliver data and/or commands to the various premises **102** in the network (or more specifically, to the devices in the premises **102** that are configured to detect such notifications). The central office **103** may also include a content server **106**. The content server **106** may be one or more computing devices that are configured to provide content to users in the homes. This content may be, for example, video on demand movies, television programs, songs, text listings, etc. The content server **106** may include software to validate user identities and entitlements, locate and retrieve requested content, encrypt the content, and initiate delivery (e.g., streaming) of the content to the requesting user and/or device.

The central office **103** may also include one or more application servers **107**. An application server **107** may be a computing device configured to offer any desired service, and may run various languages and operating systems (e.g., servlets and JSP pages running on Tomcat/MySQL, OSX, BSD, Ubuntu, Redhat, HTML5, JavaScript, AJAX and COMET). For example, an application server may be

4

responsible for collecting television program listings information and generating a data download for electronic program guide or interactive program guide listings. Another application server may be responsible for monitoring user viewing habits and collecting that information for use in selecting advertisements. Another application server may be responsible for formatting and inserting advertisements in a video stream being transmitted to the premises **102**. Another application server may be responsible for receiving user remote control commands, and processing them to provide an intelligent remote control experience.

An exemplary premises equipment may include an interface **120**. The interface **120** may, in one aspect, comprise a modem **110**, which may include transmitters and receivers used to communicate on the lines **101** and with the central office **103**. The modem **110** may be, for example, a coaxial cable modem (for coaxial cable lines **101**), a fiber interface node (for fiber optic lines **101**), or any other modem device. The modem **110** may be connected to, or be a part of, a gateway interface device **111**. The gateway interface device **111** may be a computing device that communicates with the modem **110** to allow one or more other devices in the premises to communicate with the central office **103** and other devices beyond the central office. The gateway **111** may comprise a terminal such as a set-top box (STB), digital video recorder (DVR), computer server, or any other desired computing device. The gateway **111** may also include (not shown) local network interfaces to provide communication signals to devices in the home, such as televisions **112**, additional STBs **113**, personal computers **114**, laptop computers **115**, wireless devices **116** (wireless laptops and notebooks, mobile phones, mobile televisions, personal digital assistants (PDA), etc.), and any other desired devices. Examples of the local network interfaces include Multimedia Over Coax Alliance (MoCA) interfaces, Ethernet interfaces, universal serial bus (USB) interfaces, wireless interfaces (e.g., IEEE 802.11), Bluetooth interfaces, and others.

In some embodiments, communications between the various components and devices shown in FIG. 1 may occur. For example, communications may occur in a downstream direction, an upstream direction, or in both downstream and upstream directions (e.g., via lines **101**). In some embodiments, the communications may be based at least in part on contention-based protocols, such as ALOHA communications.

FIG. 2 illustrates general hardware elements that can be used to implement any of the various computing devices discussed above. The computing device **200** may include one or more processors **201**, which may execute instructions of a computer program to perform any of the features described herein. The instructions may be stored in any type of computer-readable medium or memory, to configure the operation of the processor **201**. For example, instructions may be stored in a read-only memory (ROM) **202**, random access memory (RAM) **203**, removable media **204**, such as a Universal Serial Bus (USB) drive, compact disk (CD) or digital versatile disk (DVD), or any other desired electronic storage medium. Instructions may also be stored in an attached (or internal) hard drive **205**. The computing device **200** may include one or more output devices, such as a display **206** (or an external television), and may include one or more output device controllers **207**, such as a video processor. There may also be one or more user input devices **208**, such as a remote control, keyboard, mouse, touch screen, microphone, etc. The computing device **200** may also include one or more network interfaces, such as input/output circuits **209** (such as a network card) to communicate

5

with an external network **210**. The network interface may be a wired interface, wireless interface, or a combination of the two. In some embodiments, the interface **209** may include a modem (e.g., a cable modem), and network **210** may include the communication lines **101** discussed above, the external network **109**, an in-home network, a provider's wireless, coaxial, fiber, or hybrid fiber/coaxial distribution system (e.g., a DOCSIS network), or any other desired network.

Certain aspects of this disclosure may be implemented in connection with one or more wired or wireless networks. Aspects of this disclosure may be applied in connection with data and communication networks and protocols.

In some embodiments, a contention-based protocol may be used. For example, multiple transmitters may transmit on a given channel or frequency. A transmitter may determine whether there is energy present in the channel or frequency, and if so, the transmitter may wait for the energy to dissipate. Once the transmitter determines that there is no energy present in the channel, it may begin transmitting data. During the transmission, the transmitter may continue observing or listening to the channel or frequency to determine if a collision or contention occurs. If a collision or contention occurs, the transmitter may engage in a back-off period (which may be a random value in some embodiments) and retransmit the data upon the expiration of the back-off period.

In some embodiments, ALOHA communications or slotted ALOHA may be used. Slotted ALOHA subdivides the channel or frequency associated with standard ALOHA into multiple time slots. One or more transmitters may be assigned to particular time slots, thereby reducing the pool of transmitters that may potentially vie for the channel at a given point in time. As a result, the probability of collisions or contention using slotted ALOHA may be less than that experienced using standard ALOHA.

In theory, a particular network may operate at a maximum capacity of approximately 18%, for example, operating under an assumption of a distribution (e.g., a Poisson distribution) of messages being generated and sent in the network. Practical considerations may dictate that such a distribution is not actually present. For example, in the context of an interactive environment, there may be a spike or increase in the number of messages transmitted at various points in time. For example, in the context of an interactive program guide (IPG), if an advertisement is presented that invites users to select the advertisement (or a button or other selection mechanism associated therewith) to purchase a product or service (or engage in some other activity, such as providing enrollment information to receive more information on the product or service, or a discount regarding the same), there may be a large or significant increase in the number of messages (e.g., upstream messages) that are transmitted in the network over the period of time (or a portion thereof) that the advertisement is enabled for selection by users.

If not addressed, the periods of time in which a large number of messages are generated and sent in the network may have an impact on the performance of the network. If the number of users selecting the advertisement exceeds a threshold associated with upstream transmission resources (e.g., resources associated with communications in a direction from client devices towards a provider's device (e.g., a server)), then the indication of the particular's user's selection of the advertisement to the provider's device may be significantly delayed (e.g., beyond a threshold representative of a tolerable delay). Similarly, resources associated with downstream communications (e.g., resources associated

6

with communications in a direction from a provider's device towards one or more client or user devices) may be overwhelmed or exhausted, e.g., in response to the large number of users selecting the advertisement within a relatively short period of time. If not addressed, the consumption of resources may result in a service being unavailable or operating in a degraded state, which may in turn result in user dissatisfaction.

In some embodiments, a monitor may be used to monitor and manage resource consumption. In some embodiments, the monitor may observe traffic, messages, or data flow in real-time. In some embodiments, the monitor may observe traffic in an upstream and/or a downstream direction. In some embodiments, the monitor may comprise one or more sniffers. In some embodiments, the monitor may communicate with a network coordinator to control (e.g., throttle back) traffic.

FIG. 3 illustrates, in block diagram form, an architecture that may be used to implement one or more aspects of this disclosure. In some embodiments, the entities shown in FIG. 3 may be included in, or associated with, one or more of the entities described above in connection with FIGS. 1-2.

As shown in FIG. 3, a number of transmitters (TXs) **302a-302c** may transmit data (e.g., data **306a-306c**, respectively) to a receiver (RX) **314** via one or more networks **310**. One or more of TXs **302** may be included in one or more content creator devices, content provider devices, client or user devices or terminals, such as a set-top box, a mobile phone, a smart phone, a remote control, a display device, a computer, etc. One or more of TXs **302** may be associated with a service group. RX **314** may be included in one or more devices, such as a server. The data transmission from one or more of TXs **302** to RX **314** may occur in an upstream direction.

In some embodiments, the data **306** to be transmitted from a TX **302** to RX **314** may take place at a radio frequency (RF). The data may include a payload, such as an internet protocol (IP) payload. The data may be modulated (e.g., via the modulator shown in FIG. 3) in accordance with one or more schemes or techniques. For example, quadrature phase shift keyed (QPSK) modulation may be used in some embodiments. In some embodiments, the data may be encrypted (e.g., via the encryptor shown in FIG. 3) using one or more encryption schemes or techniques.

RX **314** may be configured to receive one or more signals at one or more frequencies. If the signal is subject to modulation or encryption, RX **314** may demodulate (e.g., via the demodulator shown in FIG. 3) or decrypt (e.g., via the decryptor shown in FIG. 3) the received signal, as appropriate. In some embodiments, RX **314** may strip off a header (e.g., via the header stripper function shown in FIG. 3), such as an Asynchronous Transfer Mode (ATM) header, to obtain the payload. A measurement of traffic (e.g., network traffic) may be based on the payload, such as an observation of the payload, in some embodiments.

A monitor may be established at point A **322** to measure and monitor the traffic (e.g., the number of messages) being sent by the TXs **302**. If the measured traffic at point A **322** indicates that the traffic exceeds a threshold amount, the monitor (or some other device associated with the monitor) may signal an entity responsible for the application that generated the traffic to throttle or reduce the traffic. For example, if RX **314** is associated with (e.g., included in) a server **344** that is responsible for providing an application, and the messages transmitted by TXs **302** are generated in response to that application, the monitor may signal the

server **344** to throttle or reduce the traffic. Exemplary techniques for throttling or reducing traffic are described below.

If the application in question relates to an advertisement, and the data **306** comprises messages that indicate user selections of the advertisement, if the monitor at point A **322** indicates the number of messages in response to the advertisement exceeds a threshold amount, the monitor may command or direct a device, such as a server or any other device upstream from the TXs **302** responsible for the content of the advertisement, to reduce or throttle the activity associated with the advertisement. In some embodiments, the threshold amount may be a function of an application in question, network capacity (such as total network capacity), etc. The command or directive may specify the action to be taken by the device to reduce or throttle the activity associated with the advertisement. Alternatively, or additionally, the command or directive may leave it up to the device responsible for the content of the advertisement to determine one or more actions to take (if any) in order to reduce or throttle the activity associated with the advertisement.

The device responsible for the tracking and/or content of the advertisement (e.g., a server) may take any number of actions in response to the command/directive received from the monitor. For example, the device may disable the application, such that users are no longer presented with an option to select the advertisement. Alternatively, or additionally, the device may stagger the application, such as by limiting the number of client devices that receive the application, e.g., (only) a subset of client devices within a service group may receive the application. The staggering may be based on time, e.g., the application may be active for only a subset of time at a first client device relative to the length of the advertisement, and the time that the application is active at the first client device may be different from the time that the application is active at a second client device. In some embodiments, the device (e.g., the server) may impose a restriction on the number of users that are able to respond to, or interact with, the application, which may result in a reduction of traffic, e.g., upstream traffic.

The input or stimulus that is responsible for generating the messages (e.g., the interactive application associated with the advertisement) may be reduced or scaled back, which may result in a corresponding reduction or throttling of the messages from TXs **302** to RX **314**. More generally, a reduction or throttling of the stimulus provided to client devices in the downstream direction may result in a reduction or throttling of traffic originating from the client devices in the upstream direction. In some embodiments, traffic may be monitored in an upstream direction, and responsive to determining that the monitored upstream traffic exceeds one or more thresholds, a controlling or throttling of traffic in a downstream direction may occur. In some embodiments, downstream data traffic may comprise portions that may require corresponding data traffic in the upstream direction.

In addition to, or as an alternative to, monitoring at point A **322**, the monitor may be located at point B **328**. For example, the monitor at point B **328** may monitor within a network, such as an RF network. The monitor may comprise a radio and may listen for on-going traffic (e.g., messaging). Locating the monitor at point B **328** may be used to determine network congestion based upon actual collisions. A more granular view of the network traffic may be obtained via a monitor located at point B **328** relative to the view obtained via a monitor located at point A **322**. Furthermore, locating a monitor at point B **328** may enable a throttling of

traffic at the instant collisions start to occur, or at some other point in time that may be correlated to, or associated with, a threshold value.

In some embodiments, the monitor may be located at a point different from (or in addition to) one or both of points A **322** and B **328** shown in FIG. 3. For example, in FIG. 4, four receivers (RXs) **314a-314d** are shown. In some embodiments, more or fewer than four RXs may be included. One or more of RXs **314a-314d** may be associated with one or more applications. One or more of receivers (RXs) **314a-314d** may include devices similar to those shown in RX **314** of FIG. 3.

Also shown in FIG. 4 is a monitor located at point C **404**. When a monitor is located at point C **404**, the monitor may aggregate traffic measurements from the RXs **314a-314d** and may take one or more actions similar to those described above in connection with a monitor located at points A **322** and B **328** shown in FIG. 3. In this manner, the architecture of FIG. 4 may facilitate resource management across a plurality of applications for a given network or set of networks.

FIG. 5 illustrates a monitor **504** in accordance with one or more aspects of this disclosure. In some embodiments, monitor **504** may be located at one or more of points **322**, **328**, and **404** described above in connection with FIGS. 3-4.

Monitor **504** may comprise a processor **510** (e.g., processor **201** of FIG. 2) and a memory **516** (e.g., ROM **202** of FIG. 2). Memory **516** may store instructions that, when executed by processor **510**, cause monitor **504** (or components or devices included in monitor **504**) to perform one or more acts or processes.

Monitor **504** may comprise a traffic measurer **522**. Traffic measurer **522** may receive as input traffic (e.g., data or messages) that appear on a communication channel or medium (e.g., lines **101** of FIG. 1). Traffic measurer **522** may store the input traffic. Traffic measurer **522** may store statistics regarding the input traffic, such as an identification of a source(s) or destination(s) of the traffic, the volume of traffic (e.g., a number of messages received per unit time), etc.

Monitor **504** may comprise a comparison monitor **528**. Comparison monitor **528** may compare the traffic measured by traffic measurer **522**, or other characteristics, against one or more thresholds. The one or more thresholds may be stored in memory **516** and may be retrieved by, or accessed at, comparison monitor **528**.

Comparison monitor **528**, which may comprise hardware, firmware, and/or software, may provide an indication of the comparison of the measured traffic to the one or more thresholds to a hardware, firmware, and/or software device, such as a director **534**. Director **534** may issue one or more commands or directives in response to the indication from comparison monitor **528**. For example, if the indication provided by comparison monitor **528** indicates that the measured traffic exceeds one or more thresholds, director **534** may issue a command or directive that causes a throttling or reduction of the traffic in accordance with the techniques described herein.

In some embodiments, a monitor (e.g., monitor **504**) may have established rules related to classes of services. Certain types of traffic may be prioritized over others. For example, messaging related to specific services such as Video On Demand (VOD) may have a priority level assigned which is higher than that of an advertising-based service.

In some embodiments, certain forms of messaging interaction may be known as "short" while others may be known as "long" form interaction. "Short" messaging may include

a limited set of interactions associated with an exchange, and long may be related to a more interactive session. An example of these might be the contrast between a Caller ID notification and an Interactive Gaming Session. The Caller ID notification might be classified as “short” and the Interactive Gaming Session might be classified as “long”. In some embodiments, a classification may be used to determine whether to request, e.g., an application server to throttle messaging ahead of the actual load on the network being realized.

In some embodiments, a taxonomy may be created for the classification of data networking exchanges used in a network (e.g., a cable network). In some embodiments, one or more upstream/downstream message exchanges may be classified according to type. In some embodiments, rules associated with a type and/or provisioned prioritizations may be applied. The applied rules may result in a throttling of lower priority or more highly interactive sessions.

The architectures and hardware/firmware/software environments described above are illustrative. One or more of the entities shown might not be included in some embodiments. For example, in some embodiments, additional entities not shown in FIGS. 1-5 may be included. In some embodiments, the entities may be arranged in a manner different from what is described.

FIG. 6 provides an example method including steps that may implement one or more aspects of this disclosure. The method may be operative in one or more environments, and on one or more devices or entities, such as the environments and devices described above in connection with FIGS. 1-5.

As shown in FIG. 6, in step 602 an application or piece of content may be provided to one or more user or client devices. The application may be associated with one or more content assets, such as a television program, an advertisement, an audio track, a text file, an email, an instant message, etc. For example, the application may provide for an interactive function or capability that is associated with the content asset(s). In some embodiments, the provided application may adhere to the Enhanced TV Binary Interchange Format (EBIF).

In step 608, traffic may be measured (e.g., by traffic measurer 522 of FIG. 5). The measured traffic may correspond to traffic generated by the one or more client devices in an upstream direction. The measured traffic may comprise data and/or messages, and may include or be generated based at least in part on the application provided in connection with step 602. For example, if the application provides an interactive selection mechanism associated with an advertisement, the traffic may signify or indicate user selections of the advertisement (or the mechanism associated therewith).

In step 614, the measured traffic may be monitored. For example, the measured traffic of step 608 may be compared (e.g., by comparison monitor 528 of FIG. 5) with one or more threshold values in step 614 to determine whether a resource (e.g., bandwidth) is being exploited in an efficient or preferential way. In some embodiments, the one or more threshold values may be determined by a network, service, or content provider. In some embodiments, the one or more threshold values may be configurable by one or more users. In some embodiments, the one or more threshold values may be adaptive over time, possibly in response to one or more events. For example, the threshold(s) may be lowered in response to recent application server failures in order to protect against additional failures.

In step 620, a determination may be made (e.g., by comparison monitor 528 of FIG. 5) whether the measured

traffic is acceptable. For example, if the monitoring of step 614 indicates that the measured traffic does not exceed a threshold (e.g., the “YES” path is taken out of step 620), then operational flow may proceed back to step 608 to continue measuring the traffic. In this manner, a loop may be maintained so long as the measured traffic is acceptable (e.g., so long as the measured traffic is within threshold limits). If, on the other hand, the monitoring associated with step 614 indicates that the measured traffic exceeds a threshold (e.g., the “NO” path is taken out of step 620), then flow may proceed to step 626.

In step 626, one or more actions may be taken to control, for example, throttle, the traffic. In some embodiments, one or more commands or directives may be issued (e.g., by director 534 of FIG. 5) to throttle the traffic. In some embodiments, the throttling of the traffic may comprise throttling, reducing, or scaling back an input or stimulus in a downstream direction, which may result in a throttling of the measured traffic in, e.g., the upstream direction. For example, a message may be sent to a server responsible for providing an application to suspend sending data associated with the application to client devices. In some embodiments, a message may be sent to a client device to disable selection or interaction with the application at the client device. In some embodiments, an indication of the throttling back may be provided to one or more users. For example, in the context of an advertisement, a display screen may be generated that advises a user that the user’s request (with respect to the advertisement) cannot be processed at this time.

In some embodiments, the throttling action may subside or expire. For example, the throttling action may expire after a predetermined amount of time has lapsed, in response to an event (e.g., upstream traffic falling below a threshold associated with step 620), etc. A user selection or interaction with an application during the time when traffic is throttled may be processed after the throttling action has expired. For example, a user device or a client device may store or buffer a user selection until a point in time is reached where the condition(s) that caused the throttling to take place have subsided.

The method described above in connection with FIG. 6 is illustrative. In some embodiments, some of the steps (or portions thereof) may be optional. In some embodiments, additional steps not shown may be included. In some embodiments, the order or sequence of the steps may be different from what is shown in FIG. 6.

Aspects of this disclosure may readily be applied to, and adapted to be operative on, one or more communication systems, including computer networks, television networks, satellite networks, telephone and cellular networks, and the like.

Although not required, functions described herein may be embodied as a method, a data processing system, and/or as a non-transitory computer-readable medium storing executable instructions. Aspects of the disclosure may take the form of an entirely hardware embodiment, an entirely software embodiment, an entirely firmware embodiment, or an embodiment combining software, firmware and hardware. The functionality may be resident in a single computing device, or may be distributed across multiple computing devices/platforms, the multiple computing devices/platforms optionally being connected to one another via one or more networks. Moreover, the structural components described herein may be distributed amongst one or more devices, optionally within a common housing or casing.

Although communications and contention based communication protocols have been discussed in the example

11

embodiments, the disclosure relates to and encompasses other communication techniques and protocols. In some embodiments, one or more network, service, or content providers (or devices or equipment associated therewith) may be configured to provide data (e.g., content) using one or more transmission techniques, and may manage traffic generated in response to the data. Moreover, while aspects of the disclosure have been described in terms of applications and advertisements, aspects of the disclosure also include other forms and types of data (e.g., text, audio, computer programs, movies, programs/shows, caller identifications, etc.) delivered or distributed via one or more schemes or techniques (e.g., video on demand (VOD)).

In some embodiments, an active video network (AVN) may be used in connection with a browser, where a client device may be provided with a content-rich (e.g., graphic rich) navigation interface. Intelligence may reside in one or more network devices, and user selections may be conveyed from client devices to the one or more network devices to operate on the user selections. More generally, aspects of disclosure may be applied to any context, e.g., environments incorporating upstream communications.

The methodological acts and processes may be tied to particular machines or apparatuses. For example, as described herein, an application may be distributed to a client or user device via one or more computing devices (e.g., servers) and that application may be rendered at the client device (e.g., at a display device associated with the client device). The application may support user interaction, and traffic generated in response to that user interaction may be measured and monitored in connection with a resource utilization. More generally, one or more computers may include one or more processors and memory storing instructions that, when executed, cause the one or more computers to perform the methodological acts and processes described herein. Furthermore, the methodological acts and processes described herein may perform a variety of functions including transforming an article (e.g., upstream traffic) into a different state or thing (e.g., throttled upstream traffic).

Aspects of the disclosure have been described in terms of illustrative embodiments thereof. While illustrative systems and methods as described herein embodying various aspects of the present disclosure are shown, it will be understood by those skilled in the art, that the disclosure is not limited to these embodiments. Modifications may be made by those skilled in the art, particularly in light of the foregoing teachings. For example, each of the features of the aforementioned illustrative examples may be utilized alone or in combination or subcombination with elements of the other examples. For example, any of the above described systems and methods or parts thereof may be combined with the other methods and systems or parts thereof described above in any order. It will also be appreciated and understood that modifications may be made without departing from the true spirit and scope of the present disclosure. The description is thus to be regarded as illustrative instead of restrictive on the present disclosure.

The invention claimed is:

1. An apparatus, comprising:

one or more processors; and

memory storing instructions that, when executed by the one or more processors, configure the apparatus to:

receive a plurality of messages, each of the plurality of messages indicating an action performed by one or more first computing devices responsive to a first advertisement sent by one or more second computing devices; and

12

send, to the one or more second computing devices and based on determining that a rate at which the messages were received satisfies a threshold, a command that causes the one or more second computing devices to change a rate at which the one or more second computing devices sends a plurality of second advertisements.

2. The apparatus of claim 1, wherein the each of the messages indicates one or more of: an enrollment action, a purchase action, or a request for webpage content.

3. The apparatus of claim 1, wherein the plurality of messages comprises indications of user selections associated with the first advertisement, and wherein the command causes the one or more second computing devices to stop sending the plurality of second advertisements for a period of time.

4. The apparatus of claim 1, wherein the instructions, when executed by the one or more processors, further configure the apparatus to:

after sending the command, receive, from the one or more first computing devices, an additional plurality of messages indicating actions performed by the one or more first computing devices responsive to the plurality of second advertisements; and

based on determining that an additional rate at which the additional plurality of messages was received satisfies an additional threshold, send, to the one or more second computing devices, another command that causes an increase in the rate of sending of the plurality of second advertisements.

5. The apparatus of claim 1, wherein the one or more first computing devices comprises a plurality of first computing devices, and wherein the command causes:

reducing, by the one or more second computing devices, the rate at which the one or more second computing devices send the plurality of second advertisements to the plurality of first computing devices by at least:

sending the plurality of second advertisements only to a first subset of the plurality of first computing devices during a first time period, and

sending the plurality of second advertisements only to a second subset of the plurality of first computing devices during a second time period.

6. The apparatus of claim 1, wherein the command causes: based on determining that the rate at which the plurality of messages was received satisfies the threshold, presenting a selectable option to remove interactive features from the plurality of second advertisements prior to sending the plurality of second advertisements.

7. The apparatus of claim 1, wherein the command causes the one or more second computing devices to adjust a frequency of interactive features in the plurality of second advertisements.

8. The apparatus of claim 1, wherein the command causes: determining a classification for the plurality of messages, and wherein determining that the rate at which the plurality of messages was received satisfies the threshold is based on the classification.

9. One or more non-transitory computer-readable media storing instructions that, when executed, cause:

receiving a plurality of messages, each of the plurality of messages indicating an action performed by one or more first computing devices responsive to a first advertisement sent by one or more second computing devices; and

sending, to the one or more second computing devices and based on determining that a rate at which the messages

13

were received satisfies a threshold, a command that causes the one or more second computing devices to change a rate at which the one or more second computing devices sends a plurality of second advertisements.

10. The one or more non-transitory computer-readable media of claim 9, wherein the each of the messages indicates one or more of: an enrollment action, a purchase action, or a request for webpage content.

11. The one or more non-transitory computer-readable media of claim 9, wherein the plurality of messages comprises indications of user selections associated with the first advertisement, and wherein the command causes the one or more second computing devices to stop sending the plurality of second advertisements for a period of time.

12. The one or more non-transitory computer-readable media of claim 9, wherein the instructions, when executed, further cause:

after sending the command, receiving, from the one or more first computing devices, an additional plurality of messages indicating actions performed by the one or more first computing devices responsive to the plurality of second advertisements; and

based on determining that an additional rate at which the additional plurality of messages was received satisfies an additional threshold, sending, to the one or more second computing devices, another command that causes an increase in the rate of sending of the plurality of second advertisements.

13. The one or more non-transitory computer-readable media of claim 9, wherein the one or more first computing devices comprises a plurality of first computing devices, and wherein the command causes:

reducing, by the one or more second computing devices, the rate at which the one or more second computing devices send the plurality of second advertisements to the plurality of first computing devices by at least:

sending the plurality of second advertisements only to a first subset of the plurality of first computing devices during a first time period, and

sending the plurality of second advertisements only to a second subset of the plurality of first computing devices during a second time period.

14. The one or more non-transitory computer-readable media of claim 9, wherein the instructions, when executed, further cause:

based on determining that the rate at which the plurality of messages was received satisfies the threshold, presenting a selectable option to remove interactive features from the plurality of second advertisements prior to sending the plurality of second advertisements.

15. The one or more non-transitory computer-readable media of claim 9, wherein the command causes the one or more second computing devices to adjust a frequency of interactive features in the plurality of second advertisements.

16. The one or more non-transitory computer-readable media of claim 9, wherein the instructions, when executed, further cause:

determining a classification for the plurality of messages, and wherein determining that the rate at which the plurality of messages was received satisfies the threshold is based on the classification.

17. A system comprising:

a first computing device; and

a second computing device;

wherein the first computing device comprises:

14

one or more first processors; and

memory storing first instructions that, when executed by the one or more first processors of the first computing device, cause the first computing device to:

receive a plurality of messages, each of the plurality of messages indicating an action performed by one or more third computing devices responsive to a first advertisement sent by the second computing device; and

send, to the second computing device and based on determining that a rate at which the messages were received satisfies a threshold, a command that causes the second computing device to change a rate at which the second computing device sends a plurality of second advertisements; and

wherein the second computing device comprises:

one or more second processors; and

memory storing second instructions that, when executed by the one or more second processors of the second computing device, cause the second computing device to:

send the first advertisement;

receive the command from the first computing device; and

change, based on the command, the rate at which the second computing device sends ones of the plurality of second advertisements.

18. The system of claim 17, wherein the each of the messages indicates one or more of: an enrollment action, a purchase action, or a request for webpage content.

19. The system of claim 17, wherein the plurality of messages comprises indications of user selections associated with the first advertisement, and wherein the command causes the second computing device to stop sending the plurality of second advertisements for a period of time.

20. The system of claim 17, wherein the instructions, when executed by the one or more first processors, further configure the first computing device to:

after sending the command, receive, from the one or more third computing devices, an additional plurality of messages indicating actions performed by one or more third computing devices responsive to the plurality of second advertisements; and

based on determining that an additional rate at which the additional plurality of messages was received satisfies an additional threshold, send, to the second computing device, another command that causes an increase in the rate of sending of the plurality of second advertisements.

21. The system of claim 17, wherein the one or more third computing devices comprises a plurality of third computing devices, and wherein the command causes:

reducing, by the second computing device, the rate at which the second computing device sends the plurality of second advertisements to the plurality of third computing devices by at least:

sending the plurality of second advertisements only to a first subset of the plurality of third computing devices during a first time period, and

sending the plurality of second advertisements only to a second subset of the plurality of third computing devices during a second time period.

22. The system of claim 17, wherein the command causes: based on determining that the rate at which the plurality of messages was received satisfies the threshold, presenting a selectable option to remove interactive fea-

15

tures from the plurality of second advertisements prior to sending the plurality of second advertisements.

23. The system of claim **17**, wherein the command causes the second computing device to adjust a frequency of interactive features in the plurality of second advertise- 5 ments.

24. The system of claim **17**, wherein the command causes: determining a classification for the plurality of messages, and wherein determining that the rate at which the plurality of messages was received satisfies the thresh- 10 old is based on the classification.

* * * * *

16