

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388742
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Keeler; Matthew et al.

Cross datacenter communication using a mesh gateway

Abstract

A method for cross datacenter service-to-service communication over a shortest network route using mesh gateways. A mesh gateway receives a protocol message from a first service directed to a second service, identifies destination information associated with the second service, recognizes a routing configuration for routing the protocol message, determines a network route based on the routing configuration, and forwards the protocol message to another gateway associated with the second service over the determined network route.

Inventors: Keeler; Matthew (San Francisco, CA), Banks; Paul (San Francisco, CA), Hashimoto; Mitchell (Manhattan Beach, CA), Dadgar; Armon (Seattle, WA)

Applicant: HashiCorp (San Francisco, CA)

Family ID: 82022680

Assignee: HashiCorp (San Francisco, CA)

Appl. No.: 18/116745

Filed: March 02, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230208746 A1	Jun. 29, 2023

Related U.S. Application Data

continuation parent-doc US 17125746 20201217 US 11601358 child-doc US 18116745

Publication Classification

Int. Cl.: H04L45/122 (20220101); H04L45/12 (20220101); H04L45/74 (20220101)

U.S. Cl.:

CPC **H04L45/122** (20130101); **H04L45/126** (20130101); **H04L45/74** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7769886	12/2009	Naseh	709/238	H04L 67/1001
8165121	12/2011	Singh	370/251	H04L 45/28
8327017	12/2011	Trost	709/225	H04L 45/02
10243834	12/2018	Shekhar	N/A	H04L 45/64
10567288	12/2019	Mutnuru	N/A	H04L 41/40
10992566	12/2020	Dave	N/A	N/A
11201897	12/2020	Miseiko	N/A	N/A
11601358	12/2022	Keeler	N/A	H04L 45/122
2007/0110063	12/2006	Tang	370/432	H04L 45/02
2010/0208662	12/2009	Fuste Vilella	370/328	H04L 45/02
2014/0064275	12/2013	Hammer	370/392	H04L 45/74
2016/0337264	12/2015	Murashko	N/A	H04L 67/63
2018/0006931	12/2017	Ellis	N/A	H04L 43/0852
2019/0034298	12/2018	Maskalik et al.	N/A	N/A
2019/0331727	12/2018	Benoit	N/A	G01R 31/086
2020/0028758	12/2019	Tollet	N/A	H04L 45/586
2020/0067734	12/2019	Hira et al.	N/A	N/A
2020/0067982	12/2019	Leach et al.	N/A	N/A
2020/0104161	12/2019	Kapur et al.	N/A	N/A
2020/0183905	12/2019	Wang et al.	N/A	N/A
2020/0259892	12/2019	Magnezi	N/A	H04L 67/1006
2021/0014113	12/2020	Guim Bernat et al.	N/A	N/A
2021/0124758	12/2020	Bottari et al.	N/A	N/A
2021/0184983	12/2020	Ramaswamy et al.	N/A	N/A
2021/0297925	12/2020	Berzin et al.	N/A	N/A
2022/0345386	12/2021	Madanapalli	N/A	A63F 13/67

OTHER PUBLICATIONS

U.S. Appl. No. 17/125,746, “Non-Final Office Action”, Apr. 13, 2022, 14 pages. cited by applicant
U.S. Appl. No. 17/125,746, “Notice of Allowance”, Oct. 31, 2022, 9 pages. cited by applicant

Primary Examiner: Sweet; Lonnie V

Attorney, Agent or Firm: FisherBroyles LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 17/125,746, filed on Dec. 17, 2020, entitled “CROSS DATACENTER COMMUNICATION USING A MESH GATEWAY”, the contents of which is incorporated by reference herein in its entirety.

TECHNICAL FIELD

(1) This disclosure generally relates to techniques for cross datacenter communication between web services using mesh gateways. Specifically, the present disclosure provides a technique to configure gateways with optimal routes to target datacenters in order to enable cross datacenter communication.

BACKGROUND OF THE INVENTION

(2) A software application may span multiple web-services within different data centers. Securely connecting these web-services associated with multiple run-time platforms or environments involves complex and costly configurations between the services. Spanning web-services within different clouds, a cloud within datacenter, and connections between a modern scheduler and virtual machines also require a careful network and security configuration.

(3) Typically, web-services within different environments are connected using a point to point VPN connection. VPN uses a public network infrastructure, by means of tunneling technology to achieve data security transmission of similar private network. Correctly configuring and operating VPN tunnels to establish point to point VPN connections within multiple platforms or environments is a costly and complex process.

(4) While a technique that allows service discovery across datacenter boundaries by differentiating between WAN and LAN addresses of a service and automatically using the WAN addresses for a direct connection between web-services exists, this technique does not ensure that subnets between two environments do not overlap. Additionally, this approach also requires established VPN set-up between environments.

(5) Typically, after the VPN connection is established, a host within a local gateway within a datacenter connects to a host within a remote gateway through the secure VPN tunnel. A connection consists of the IP addresses and host names of the two gateways, the subnets behind them, and the shared secrets for the two gateways to authenticate to each other. When thousands of users access applications, defining policies and synchronizing them across different locations is a complex task. In addition, deploying, configuring and managing VPNs in tens or hundreds of cloud instance is expensive, time-consuming and risky.

(6) While there are numerous techniques for service-to-service communication within different environments, a simple solution for service-to-service communication using optimal network routes that solves the various networking problems and complexities that arise in multi-cloud, multi-cluster, multi-platform, hybrid-cloud environments is not available. Accordingly, there is a need for a framework that enables efficient cross datacenter service-to-service communication using an optimal network route.

BRIEF SUMMARY OF THE INVENTION

(7) The present disclosure describes techniques for cross datacenter service-to-service communication using one or more mesh gateways. Various embodiments are described herein, including methods, systems, non-transitory computer-readable storage media storing programs, code, or instructions executable by one or more processors, and the like. Embodiments of the invention add a new concept for connecting to one or more datacenters using “Mesh Gateway(s)” to solve the various networking problems and complexities that arise in multi-cloud, multi-cluster, multi-platform, hybrid-cloud environments. For the sake of brevity any mention of “gateway” later

in this disclosure refers specifically to a mesh gateway.

(8) In an example embodiment, a mesh gateway may receive a protocol message directed to a service within a target datacenter. The mesh gateway may sniff into the protocol message to identify information indicating the destination address of the service and the target datacenter. The destination hostname of the service can be within a SNI header of the protocol message. In the above embodiment, the protocol message may be a part of a Transport Layer Security (TLS) Server Name Indication (SNI) handshake originated from a client service.

(9) In certain embodiments, the mesh gateway may extract from the protocol message information relating to the routing configuration for routing the protocol message to the service. The routing configuration for routing the protocol message may be single-hop routing or multi-hop routing. The mesh gateway may determine a network route for routing the protocol message to the service based on the routing configuration.

(10) In an example embodiment, if the routing configuration for routing the protocol message is configured to be single-hop mesh gateway routing, the gateway may further determine whether the gateway has a direct connection to the target datacenter. In the above embodiment, if the mesh gateway has direct connection to the target datacenter, the mesh gateway uses the direct connection as a network route to route the protocol message to the target datacenter.

(11) In an alternate embodiment, if the routing configuration for routing the protocol message is set to single-hop routing and there is no direct connection to the target datacenter, the mesh gateway may determine a distance(s) to one or more next-hop datacenter(s) directly accessible by the mesh gateway. The mesh gateway may also determine distance(s) to the target datacenter as advertised by the next-hop datacenter(s).

(12) In the above embodiment, the mesh gateway may calculate total distance(s) to the target datacenter based on the distance(s) to the next-hop datacenter(s) and the distance(s) as advertised. The mesh gateway may further use a network route with a shortest total distance to the target datacenter, selected from the available network routes to the target datacenter.

(13) In certain embodiments, if the routing configuration configured for routing the protocol message is multi-hop routing, the gateway may determine distance(s) to next-hop datacenter(s) from the gateway. The distance(s) to the next-hop datacenter(s) may be determined using network coordinates of the datacenter associated with the first gateway device and the next-hop datacenter(s).

(14) In the above embodiment, the mesh gateway may further determine distance(s) to the target datacenter as advertised by next-hop datacenter(s). The mesh gateway may further calculate total distance(s) to the target datacenter based on the distance(s) to the next-hop datacenter(s) and the distance(s) as advertised. In the above embodiment, the gateway device identifies a network route with a shortest total distance to the target datacenter from the calculated total distance(s).

(15) In an example embodiment, the mesh gateway may filter a reference to the mesh gateway (itself) from the network route for routing the protocol message to the target datacenter. Filtering the gateway from the network route may prevent traffic to be routed using cyclical routes. In certain embodiments, after network route has been determined, the mesh gateway forwards the protocol message to another gateway device within the target datacenter by transmitting the protocol message through the determined network route.

(16) In the above embodiment, the gateway may use at least one of a Local Area Network (LAN) address and a Wide Area Network (WAN) address to route the protocol message to the service within the target datacenter. In the above embodiment, the gateway may use either layer 4 or layer 7 load balancing protocols for routing the protocol message. These illustrative embodiments are mentioned not to limit or define the disclosure, but to provide examples to aid understanding thereof. Additional embodiments are discussed in the Detailed Description, and further description is provided there.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Features, embodiments, and advantages of the present disclosure are better understood when the following Detailed Description is read with reference to the accompanying drawings.
- (2) FIG. 1 depicts an example of a conventional cross datacenter service-to-service communication using a VPN gateway.
- (3) FIG. 2A depicts an example of a cross datacenter service-to-service communication using a mesh gateway and WAN Gossip in accordance with certain embodiments.
- (4) FIG. 2B depicts an example of a block diagram showing interactions between components within the first and second datacenters for cross datacenter communication.
- (5) FIG. 3 depicts an example of a cross datacenter service-to-service communication using mesh gateway routing without using network area connections in accordance with certain embodiments.
- (6) FIG. 4 depicts an example of cross datacenter service-to-service communication pairwise groups of datacenter connections.
- (7) FIG. 5 depicts an example of cross datacenter service-to-service communication via mesh gateway using routes configured for the mesh gateway.
- (8) FIG. 6 depicts an example of cross datacenter service-to-service communication via multi-hop mesh gateways between three datacenters.
- (9) FIG. 7 depicts an example of cross datacenter service-to-service communication via multi-hop mesh gateways between four datacenters.
- (10) FIG. 8 depicts an example of steps performed to compute network route for cross datacenter service-to-service communication via multi-hop mesh gateways.

DETAILED DESCRIPTION OF THE INVENTION

- (11) FIG. 1 depicts a conventional method of network communication that includes end users **110**, client applications **120**, cloud management functionality **130**, internet **140**, VPN gateway **150**, data center **160**, web service **162**, service controller **164**, data server **166**, data center **170**, web service **172**, service controller **174**, and data server **176**. In an example embodiment, end users **110** may connect to one or more client applications **120**. End users may access one or more client applications **120** via computing devices such as a smartphones, laptops, desktop computers, tablets, etc.
- (12) In an example embodiment, one or more client applications **120** consists of software and/or hardware that relies on cloud computing for application delivery. In an example embodiments, a client application may be a web browser, a mobile application, or a web application. In an example embodiment, client applications **120** may issue a request to the cloud management platform **130** to access cloud computing resources.
- (13) In an example embodiment, the cloud management platform **130** may manage deployment and operations associated across multiple cloud service infrastructures (e.g., web service **162** and web service **172**). In above embodiment, the cloud management platform **130** may allow one or more client applications to use web service **162** within data center **160** and web service within data center **170** using VPN gateway **150**.
- (14) A typical VPN gateway **150** enables communication between one or more web services (e.g., service **162** and service **172**) within different datacenters or the same datacenter using a VPN infrastructure. VPN infrastructure enables secure network communication between multiple services within one data center or between multiple data centers (data center **160** and data center **170**) using VPN tunnels between services.
- (15) In an example embodiment, network traffic traveling between datacenter would be encrypted by originator service **162** and decrypted by the receiver service **172**. More specifically, Internet Protocol Security (IPsec) VPN tunnels are used for communication between datacenters. Once a

VPN connection is established, hosts behind the local gateway within datacenter **160** connect to hosts behind a remote gateway within the datacenter **170** through the secure VPN tunnel (VPN gateway **150**). In an example embodiment, a VPN gateway **150** may be built into the cloud management platform **130**.

(16) FIG. 2A depicts an example of a cross datacenter service-to-service communication using one or more mesh gateways within datacenters. In an example embodiment, when a service **162** needs to communicate with service **172** which is located in a different physical datacenter **170**, the traffic may be proxied securely through the mesh gateway **230**. In the above embodiment, the mesh gateway forwards the traffic to the correct destination service **172**.

(17) In the above embodiment, each mesh gateway (e.g., **210**, **220**, **230**, and **240**) may route to all other datacenter they are aware of that contain a mesh gateway. In the above example embodiment, a datacenter (**160**) may discover the other datacenter (**170**) through a WAN (Wide Area Network) federation. The WAN federation allows servers in different datacenters to be federated exclusively through mesh gateways. With WAN federation, all servers in first datacenter may be directly connectable over their WAN-advertised network address from servers in a second datacenter where both first and second datacenters are part of a multi-datacenter cluster.

(18) In an example embodiment, because the traffic between datacenter **160** and datacenter **170** is proxied by the mesh gateways (**230** and **240**), a web-service **162** does not need a direct connection to the destination service **172**. In an example embodiment, the web service **162** needs access only to mesh gateway **230** in order to establish communications with entities outside its datacenter **160**.

(19) In an example embodiment, the mesh gateway **230** does not have to access a destination server certificate of or keys to service **172**. In the above embodiment, gateway **230** may use Transport Layer Security (TLS) communication protocol extension Server Name Indication (SNI) to communicate with service **172**.

(20) In an example embodiment, during the TLS handshake between services **162** and **172**, service **162** may send a destination (e.g., service **172**) hostname in the SNI header that is part of the TLS handshake. In an example embodiment, gateway **230** may peek at the first message of the handshake and read the SNI header. The SNI header may describe a specific service (e.g., service **172**) that is being requested or a subset of an instance of the service (e.g., service **172**) using a pre-configured name. Accordingly, gateway **230** determines the destination without having to decrypt and re-encrypt the traffic between service **162** and service **172**.

(21) In an example embodiment, gateway **230** may further proxy the encrypted traffic from service **162** to service **172** at network transport layer L4. In the above embodiment, this approach allows service end to end TLS connection without decrypting the traffic flow even if gateway **230** is compromised.

(22) In an example embodiment, mesh gateway **230** may be an ingress and/or egress type of gateway. In the above embodiment, mesh gateway **230** may accept incoming connections (e.g., from service **172**) on a public interface and proxy them to services within its datacenter **160**. In an example embodiment, datacenter **160** includes more than one service (e.g., service **162**). Ingress type gateway **230** may direct incoming traffic to these services.

(23) In an example embodiment, gateway **230** may be an egress type gateway that accepts incoming connections from services (e.g., service **162**) within its datacenter **160** and proxies the traffic to the ingress type gateway of the target data center. In an example embodiment, as depicted in FIG. 2, if gateway **230** is an egress type gateway that accepts incoming connections from service **162** to connect to service **172**, the gateway **230** may proxy the traffic to gateway **240**, where the gateway **230** may be federated to connect to the gateway **240** within the datacenter **170** using a WAN gossip pool. The WAN gossip pool may include a list of server members from different datacenters within a cloud cluster (e.g., cluster of datacenters) where the server members of the gossip pool are interconnected. In the above embodiment, gateway **240** may be an ingress type gateway that accepts incoming connections (e.g., from service **162**) on a public interface and proxy

them to services within its datacenter **170**.

(24) In an example embodiment, egress type gateway **230** may allow cross datacenter communication without requiring sidecar proxies (e.g., sidecar proxy **271**) of service **162** to be aware of any routing of the datacenter **160** even when they may need to make requests to datacenter **170**. In the above example, egress gateway **230** may be responsible for keeping track of the external data centers and services. Additionally, egress gateway **230** may provide better load balancing because when gateway **230** is used by multiple sidecar proxies, it has more knowledge about overall traffic to better spread the load.

(25) In an example embodiment, gateway **230** may be an intermediate type gateway that combines functionalities of both ingress and egress type gateways. This type of combined gateway may accept connections from a public interface and may proxy them to other ingress gateways or other intermediate gateways within other datacenters. Intermediate gateways may exist in a federated network where all datacenters don't have a direct connection to all other datacenters.

(26) In an example embodiment, the gateway **230** may be an ingress, egress, and intermediate type gateway. In an example embodiment, the gateway **230** may be an intermediate type gateway that is used for all communications within a single datacenter **160**. In the above embodiment, within datacenter **160** each sidecar proxy within service **162** may not be required to load balance connections to services independently. Instead the sidecar proxies may use the gateway **230** to load balance connections since it has much more knowledge of overall traffic throughout the cluster. Accordingly, the gateway **230** may make better routing decision than sidecar proxies.

(27) In an example embodiment, a gateway **230** may be configured within datacenter **160**. After configuration of gateway **230**, gateway **230** may be registered and enabled in the datacenter **160**. In an example embodiment, once the gateways **230** and **240** are enabled, sidecar proxies associated with web services **162** and **172** may reconfigured automatically to send their traffic to gateways **230** and **240** in the local datacenter which will then route to gateways in the remote datacenters and remote destination services. In the above embodiment, the web service **162** and service **172** are instances of two different web services.

(28) In an example embodiment, once the gateway **230** is enabled, services within datacenter **160** are also configured to take advantage of gateway **230**. In the above embodiment, services within datacenter may have list of configured upstreams (e.g., connections to other services). In an example embodiment, one or more of upstream are configured to use the gateway **230**. In the above embodiment, the gateway **230** may modify the destination of the original outgoing request to other service. Accordingly, service proxies associated with services (e.g., service **162**) may be configured to make their TCP connection to the gateway **230** instead of the upstream connection to the final service destination.

(29) In an example embodiment, a new field “mesh gateway” may be added to the upstream configuration of the one or more web services within the datacenter. A setting in the upstream configuration for services may indicate whether a gateway should be used and the desired behavior. In the above embodiment, the upstream configuration for a service (e.g., service **162**) may be stored and managed by the sidecar proxy (e.g., sidecar proxy **271**) associated with a service (e.g., service **162**). In an example embodiment, the mesh gateway field may have one of the four values: 1) the empty string (“ ”) may represent the default gateway usage and allows overriding this setting from a global service configuration, 2) the string “none” explicitly disables using gateways for the upstream, 3) the string “remote” enables a proxy to utilize a direct connection to an ingress gateway in a remote datacenter for the upstream, and 4) the string “local” enables the proxy to utilize an egress gateway within the local datacenter for the upstream.

(30) In an example embodiment, when the gateway **230** is enabled, upstreams associated with service **162** may be configured differently. When the gateway **230** is enabled, an upstream endpoint (e.g., destination service **172**) may be set to an address of a gateway (e.g., gateway **230**) instead of an address associated with a destination service (e.g., service **172**). In an example embodiment,

depending upon the upstream's configuration, an ingress gateway in remote datacenter (e.g., 170) or an egress gateway (e.g., 230) within the local datacenter (e.g., 160) may have direct connection to the service's proxies.

(31) In an example embodiment, along with changing the upstream's target address, SNI context may be added to the upstream for the connection to be routed appropriately. For egress and intermediate type gateways, whether upstreams are configured to point at an ingress gateway in remote datacenter, proxies of egress/intermediate gateway do not require any changes. All the proxies may need to know is the next hop address to make TCP connections.

(32) In the above embodiment, a gateway (230) may be implemented using a configured envoy proxy. In an example embodiment, for egress type gateway 230, an envoy may sniff SNI header within service 162 and transmit the encrypted data to correct destination (e.g., service 172) using a TCPproxy or SNI cluster network filter (e.g., network filter that uses the SNI values in a TLS connection as the upstream cluster name). Once the destination is identified by sniffing the SNI header during the initial TLS handshake, the rest of the stream may stay encrypted during the transmission.

(33) As discussed in the above embodiment, the SNI header may be sniffed during the initial TLS handshake to retrieve destination information. For egress type gateway 230, a destination may be identified by matching domain suffixes in the TLS SNI header to those configured for other datacenters. In an example embodiment, for an egress type gateway 230, once the sniffed SNI information is matched to an external datacenter, the gateway 230 transmits the message or packet to the destination service within the external datacenter or to a next-hop gateway to transmit the packet to destination. In above embodiments, the SNI sniffing and routing mechanisms for both egress and ingress type of gateway transmissions are identical.

(34) In an example embodiment, destination information or server names within the SNI header may be specified as a fully qualified domain name having the following format.

(35) <subset>.<service>.<namespace>.<datacenter>.<internal?>.<trust domain>

(36) In an example embodiment, a destination information for a service in DC2 170 in the above format may be described as follows.

(37) v 1.web.default. dc2.int.55084bf6-fb26-4d75-b812-91f5de1c7335.consul

(38) In the above example embodiment, the <subset> name may be defined in a service-resolver configuration entry field for a cloud networking infrastructure such as Consul® by HashiCorp Inc. The service-resolver configuration entry is stored by the Consul servers (through Raft). Consul will use this configuration entry to generate the gateways configuration which will route the subset accordingly. The proxy itself is unaware of the configuration entry

(39) Additionally, the string contains a service name after the subset name. In an example embodiment, a namespace may identify the service. In the above example, the namespace is set to default. Additionally, the string may include the name of the datacenter (e.g., DC2) the service is residing within.

(40) In an example embodiment, a gateway 230 may have three types of addresses: a LAN address, a WAN address, and one or more tagged addresses. Configuring tagged addresses and the LAN/WAN address establishes how everything else talks to that gateway 230. In some implementations, a single agent can manage multiple gateways regardless of their address configuration. Gateway 230's LAN address is used by all services in DC 160 to do egress out of the DC through the Gateway. Its WAN address is used by default by all gateways residing in other datacenters when they need to send traffic into DC 160 from services in their own datacenters. The tagged addresses idea here would allow for advertising a different IP for individual datacenters to use. For example with the tagged address of dc:DC3 (element 250), then gateway 210 would use this tagged address when sending service traffic to DC1 160. This address would be ignored by gateways 240 and 220, which will still use the WAN address of gateway 230 to send traffic to DC1 160.

(41) Specifically, in an example embodiment, gateway **230** may be implemented in one or more phases. Initially, the tagged address may not be added. The gateway **230** may be contacted over a WAN address. This arrangement may allow an agent (e.g., **273**) to manage a single gateway (e.g., **230**) or multiple gateways, and force the gateway to reside on a same network namespace as the agent. In the next phase, tagged addresses may be added to the gateway **230**. Accordingly, the tagged address may now allow a single agent (e.g., **273**) within a datacenter (**DC1 160**) to manage multiple gateways (e.g., gateway **230**) and allow the gateways to use a separate network from the agent.

(42) In an example embodiment, an agent (e.g., **273**) within a data may use an xDS (Extended Data Service)/ADS (Aggregated Discovery Service) protocol to configure the gateway **230**. A local agent (e.g., **273**) within the datacenter may perform queries against a catalog to discover all mesh gateway instances in each potential target datacenter. For each of the datacenters, the gateway may be configured to route connections destined for or through that datacenter to the discovered mesh gateway instances.

(43) In an example embodiment, an ingress type gateway may be configured with one or more public listeners that has a tls inspector listener filter enabled as well as the sni cluster and tcpproxy network filters enabled. TLS Inspector listener filter allows detecting whether traffic flow (e.g., including protocol message) appears to be TLS or plaintext, and if it is TLS, it detects the Server Name Indication and/or Application-Layer Protocol Negotiation associated with the traffic. The sni cluster within the gateway **230** is a network filter that uses an SNI field in a TLS connection as the upstream cluster name (e.g., **DC2 170**). For example if there is a service “web” (e.g., service **172**) and the gateway **240** is proxying connections for “dc2” then a cluster name that starts with web.de.sup.-Fault.dc2 will be configured with load balanced endpoints corresponding to each service instance of “web” within “DC2.”

(44) The TCP proxy filter within the gateway **230** performs basic **1:1** network connection proxy between downstream clients (e.g., clients within **DC1 160**) and upstream clusters (e.g., services within **DC2 170**). The TCP proxy filter may within gateway **230** also respect the connection limits imposed by each upstream cluster. The TCP proxy filter checks with the upstream cluster's resource manager if it can create a connection without going over that cluster's maximum number of connections, if it can't the TCP proxy will not make the connection.

(45) The following code may be used to set up the listener and network filters within the gateway **230**:

```
(46) TABLE-US-00001  “Service”: “gateway-service”,  “Port”: 8443,  “Kind”: “mesh-gateway”,  “Address”: “172.16.0.1”,  “TaggedAddresses”: {    “lan”: {      “address”: “172.16.0.1”,      “port”: 8443,    },    “wan”: {      “address”: “198.18.0.1”,      “port”: 8443,    },    “dc:dc3”: {      “address”: “198.18.20.3”,      “port”: 8443,    }  } }
```

(47) In an example embodiment, egress and intermediate gateways may be configured in similar manner. Within the gateway **230**, endpoints of the target cluster (e.g., **DC2 170**) may be configured as the addresses of ingress gateways (e.g., gateway **240**) servicing the target datacenter. In an example embodiment, both the ingress and egress parts for an intermediate gateway may be configured together. In an example embodiment, gateways **230** and **240** may be compatible with L7 routing.

(48) The following code may be used to set up the egress type gateway.

```
(49) TABLE-US-00002 static resources: listeners: - name: sni automatic address: socket address: address: 1.2.3.4 port value: 443 Listener filters: - name: envoy.listener.tls inspector typed config: filter chains: - filter chain match: server names: [“*.dcl”] filters: - name: envoy.tcp proxy config: stat prefix: tcp.proxy.dcl cluster: dcl
```

(50) In an example embodiment, a Remote Procedure Call (RPC) by services within the datacenter

may be used to discover gateways in the datacenter. In an example embodiment, to discover all gateways for a datacenter, a RPC endpoint may be modified to take a specific “kind” of service (e.g., gateway service) and only return the services for that kind. The following code shows an example of configuration for a gateway service. #configuring a gateway service with this name service “gateway-name” { policy=“write” kinds=[“mesh-gateway”]}

(51) As discussed in the above example, a “kinds” variable is set to “mesh-gateway” for gateway services within the datacenter. Accordingly, all gateways within the datacenter may be discovered by finding services for which “kinds” is set to “mesh-gateway.” The service definition will be updated so the various catalog endpoints may be able to handle gateways but the APIs for definition itself would be unaltered.

(52) As shown in FIG. 2A, in an example embodiment, all server nodes in each datacenter (e.g., servers **166**, **176**, **252**, and **262**) may gossip with all the server nodes in every other data center (e.g., servers **166**, **176**, **252**, and **262**). In the above embodiment, through the use of a gossip protocol, the servers (**166**, **176**, **252**, and **262**) may determine a list of all connected servers and all connected datacenters via WAN gossiping. In an example embodiment, the gossip protocol may be a Scalable Weakly-consistent Infection-style Process Group Membership (SWIM) type protocol. The gossip protocol may provide the server (e.g., server **166**) a view of health and location of every other server across the Wan Area Network (WAN).

(53) The above embodiment, the method of WAN federation or gossip requires full WAN connectivity from every server with every other server in all datacenters. In this mode, all server nodes in each datacenter (e.g., server **166**, server **176**, etc.) may gossip with all the server nodes in every other datacenter (e.g., server **166**, server **176**, etc.).

(54) In the above embodiment, through the use of a protocol, the server(s) may determine a list of all connected servers and determine list of all connected datacenters. This method may require full WAN connectivity from every server with every other server in all datacenters. In an example embodiment, for datacenter communication between datacenters **160** and **260**, WAN connectivity between both datacenters **160** and **260** is required.

(55) FIG. 2B depicts an example of a block diagram showing interactions between components within the first and second datacenters for cross datacenter communication. Specifically, the block diagram shows cross datacenter communication between services **162** and **172**. In an example embodiment, service **162** may be an instance or subset of a web service. In an example embodiment service **172** may be an instance or subset of a different web service. In an example embodiment, as discussed, when a service **162** needs to communicate with service **172**, which is located in a different physical datacenter **170**, the traffic may be proxied securely through the mesh gateways **230** and **240**.

(56) In an example embodiment, the web service **162** may be connected to a sidecar proxy **271** and client agent **272**. The mesh gateway **230** may be connected to a different client agent **273** and server **166**. The sidecar proxy **271** and mesh gateway **230** may communicate with their client agents (e.g., **272** and **273**) using a xDS (service discovery) proxy configuration protocol. Client agents **272** and **273** may use Remote Procedure Call (RPC) to discover a global state of the cluster in order to produce the configuration for mesh gateway **230** and sidecar proxy **271**.

(57) In an example embodiment, the servers (e.g., server **166**) may issue remote procedure calls to other servers (e.g., server **176**) to query specific availability information for services (e.g., service **172**) in that datacenter (e.g., datacenter **170**), in order to configure gateway routes correctly. In an example embodiment, as shown in FIG. 2B, link **280** shows server-server communication flow using RPC calls.

(58) In an example embodiment, link **290** of FIG. 2B, shows the flow of a service request from service **162**. A service request for a target service (e.g., service **172**) may be sent from service **162** to sidecar proxy **271** of the service **162**. The service request **290** may be forwarded to the mesh gateway **230**. The gateway then forwards the request **290** to mesh gateway **240** (e.g., the gateway

associated with the datacenter of the target service) The mesh gateway **240** may communicate the request **290** to target service **172** through sidecar proxy **275**.

(59) FIG. **3** depicts an example of a cross datacenter service-to-service communication using mesh gateway routing without using network area connections in accordance with certain embodiments. As disclosed in the above embodiment describing FIG. **2A**, after the servers have determined a list of datacenters, mesh gateways within the datacenters can be configured with routes to all other datacenters. In an example embodiment, the routes to other datacenters are all direct or single-hop routes where a mesh gateway in the final destination datacenter routes the traffic (e.g., protocol messages) internally within that datacenter.

(60) In an example embodiment, as shown in FIG. **3**, four datacenters (DC1 **160**, DC2 **170**, DC3 **250**, and DC4 **260**) all have mesh gateways. Each mesh gateway within the four datacenters may be configured with three direct routes. In the above embodiment, each of the three routes is for each of the three datacenters. In an example embodiment, as shown in FIG. **3**, routes **1**, **2**, and **3** are direct routes from DC1 to DC2, DC4, and DC3. Similarly, routes **4**, **5**, and **6** are configured direct links from DC3 to DC1, DC2, and DC4. Routes **7**, **8**, and **9**, as disclosed in FIG. **3** provide direct connection from DC2 to DC4, DC3, and DC1 as well as routes **10**, **11**, and **12** provide direct connection from DC4 to DC2, DC1, and DC3.

(61) In an example embodiment, each datacenter (DC1 **160**, DC2 **170**, DC3 **250**, and DC4 **260**) may contain more than one mesh gateway. In the above embodiment, the traffic may be divided between multiple mesh gateways within a single datacenter for routing the traffic to other connected datacenters. This arrangement may further involve load balancing connections between multiple gateways within a singular datacenter. In an example embodiment, a sidecar proxy (e.g., **271** as shown in FIG. **2B**) may be configured with all peer gateways it can use within the data center (e.g., gateway **230**). The sidecar proxy may select a gateway from all of the peer gateways based on its configured load balancing algorithms such as Round Robin or Least Connections.

(62) FIG. **4** depicts an example of cross datacenter service-to-service communication using pairwise groups of datacenter connections. In an example embodiment, as shown in FIG. **4**, instead of a WAN federation, network area (e.g., Local Area Network (LAN) or Wide Area Network (WAN)) connections may be used to create pairwise groups of datacenter connections. The configuration disclosed in the above embodiment, as shown in FIG. **4**, allows only selective network reachability between the datacenters.

(63) As shown in FIG. **4**, network areas may be set up between DC1 (**160**) and DC2 (**170**), DC1 (**160**) and DC3 (**250**), DC1 (**160**) and DC4 (**260**), DC2 (**170**) and DC3 (**250**), and DC4 (**260**) and DC5 (**420**). In the above embodiment, because there is no network area configured between DC1 and DC5, these two datacenters will be unaware of each other. Accordingly, DC1 may not communicate with DC5. Similarly, other datacenters in FIG. **4** that do not have direct connection between them will be unaware of each other. However, DC1 may communicate with DC2 and DC3. In the above embodiment, accordingly, the mesh gateways deployed in each datacenter may not have routes to all datacenters but only the subset of those datacenters of which the datacenters are aware.

(64) FIG. **5** depicts an example of cross datacenter service-to-service communication via mesh gateway using routes configured for the mesh gateway. As shown in FIG. **4**, DC1 (**160**) and (65) DC2 (**170**), DC1 (**160**) and DC3 (**250**), DC1 (**160**) and DC4 (**260**), DC2 (**170**) and DC3 (**250**), and DC4 (**260**) and DC5 (**420**) may be connected using network areas. Mesh gateways within these datacenters may be configured with network area routes (LAN or WAN IP addresses) to communicate between services within connected datacenters.

(66) In an example embodiment, as shown in FIG. **5**, mesh gateway routes for DC1 include 1, 2, and 3. Similarly, a mesh gateway routes for DC2 are 4 and 5, mesh gateway route for DC3 are 6 and 7, mesh gateway routes for DC4 are 8 and 9, and mesh gateway routes for DC5 are 10. As shown in FIG. **5**, services within DC1 (**160**) may communicate to services within DC2 (**170**), DC3

(250), and DC4 (260). However, due to the lack of a route between DC3 and DC4 as well as between DC1 and DC5, services in those datacenters may not be able to communicate with each other.

(67) FIG. 6 depicts an example of cross datacenter service-to-service communication via multi-hop mesh gateways between three datacenters. As shown in FIG. 6, cross datacenter communication via mesh gateways and network areas may be extended to provide multi-hop routing between datacenters. In an example embodiment, as shown in FIG. 6, mesh gateways within datacenters may be configured with routes to connected datacenters similar to the embodiment discussed in FIG. 5.

(68) In the above embodiment, along with the direct route configuration, servers in each datacenter may advertise the other datacenters that the servers may have knowledge about in addition to their own datacenters. In the above embodiment, as shown in FIG. 5, mesh gateways may be configured with the ability to route to each known datacenter except for the route the mesh gateways may target a mesh gateway in a next-hop datacenter. In an example embodiment, next-hop datacenter may act as a point to point interface.

(69) As shown in FIG. 6, the mesh gateway 230 in DC1 160 may be configured with the ability to route the traffic to DC3 but instead of making a direct connection, it may make the connection to a mesh gateway 240 in DC2 170 which would then forward it along to DC3 250. As shown in FIG. 6, connection 1 between mesh gateway 230 and mesh gateway 240 is a direct route between DC1 and DC2. However, connection 2 between mesh gateway 240 and mesh gateway 210 is a part of multi-hop route (connections 1 and 2) from DC1 to DC3.

(70) In an another embodiment, as shown in FIG. 6, mesh gateway 210 in DC3 250 may be configured with the ability to route the traffic to DC1 but instead of making a direction connection, it may make the connection to the mesh gateway 240 in DC2 170 which would then forward it along to DC1 160. As shown in FIG. 6, connection 5 between mesh gateway 210 and mesh gateway 240 is a direct route between DC3 and DC2. However, connection 6 between mesh gateway 240 and mesh gateway 230 is a part of a multi-hop route (connections 5 and 6) from DC3 to DC1.

(71) In the above embodiment, the routing mechanism may be extended to with additional datacenters and hops to go through when routing. In an example embodiment, instead of having only one valid next-hop destination datacenter, there may be multiple next-hop destination datacenters. In an example embodiment where multiple next-hop datacenters are configured as part of a route or path to a target datacenter, a mesh gateway may ignore potential path with cycles in them and choose a best path.

(72) FIG. 7 depicts an example of cross datacenter service-to-service communication via multi-hop mesh gateways between four datacenters. As shown in FIG. 7, traffic from DC4 (260) to DC3 (260) may be routed in number of ways. In an example embodiment, it may go directly via its own direct connection 7, as disclosed in FIG. 7. In an alternative embodiment, it may be routed through DC1 (160) using direct connection 2 and multi-hop connection 9. Similarly, as shown in FIG. 7, mesh gateway 210 of DC3 (260) may directly route traffic to the mesh gateway 220 of DC4 (260) using direct connection 8, as disclosed in FIG. 7.

(73) In above examples, a mesh gateway may prefer a multi-hop connection instead of a direction connection when a direct connection is down or slow. In the above embodiment, when a direct connection 7 (as disclosed in FIG. 7) from DC4 (260) to DC3 (260) is down or slow for some reason, it may be efficient to get a protocol message routed through DC1 (160).

(74) In an alternate embodiment, a mesh gateway may always use single-hop or direct connection for traffic in all scenarios even if the direct connection is down or slow if a user of the service prefers single-hop connection and configures the mesh gateway to use only single-hop or direct connection to route the traffic.

(75) FIG. 8 depicts an example of steps performed to compute network route for cross datacenter

service-to-service communication via multi-hop mesh gateways. In an example embodiment, as disclosed in step **801**, a mesh gateway (e.g., gateway **210**) may receive a protocol message directed to a service within a target datacenter (e.g., DC4 **260**). In an example embodiment, the protocol message may be a part of a Transport Layer Security (TLS)-Server Name Indication (SNI) handshake originated from a client service.

(76) In an example embodiment, as shown in step **802**, the mesh gateway (e.g., gateway **210**) may identify the destination information associated of the service and the target datacenter (e.g., DC4 **260**) included in a SNI header of the protocol message. In the above embodiment, the mesh gateway device (e.g., gateway **210**) may sniff at the SNI header of the protocol message associated with the TLS-SNI handshake to determine the destination information for the service.

(77) In an example embodiment, as shown in step **803**, the mesh gateway (e.g., gateway **210**) may further determine a routing method configured for routing the protocol message. In the above embodiment, the routing method is configured according to the preference of a user of the service. In the above embodiment, the routing configuration for routing the protocol message is either single-hop routing (as shown in step **805**) and a multi-hop routing (as shown in step **804**). In a single-hop routing, a direct connection (e.g., connection **8** (as shown in FIG. 7)) from the mesh gateway (e.g., gateway **210**) is used to transmit the protocol message to the target datacenter (e.g., datacenter DC4 **260**). In multi-hop routing method, the mesh gateway (e.g., gateway **210**) may use a multi-hop route or a route through one or more intermediate datacenters (e.g., connections **11** and **1** as shown in FIG. 7) to transmit the protocol to the target datacenter (e.g., datacenter DC4 **260**).

(78) In the above embodiment, upon determining that the routing configuration for the mesh gateway (e.g., gateway **210**) is a single-hop routing, under step **805**, the mesh gateway (e.g., gateway **210**) may further determine whether the mesh gateway (e.g., gateway **210**) has direct connection (e.g., connection **8**, as shown in FIG. 7) to a mesh gateway (e.g., gateway **220**) of the target datacenter (e.g., datacenter DC4 **260**) under step **806**. Upon determining that the direct connection to the target datacenter (e.g., datacenter DC4 **260**) exists, the mesh gateway may use the direct connection (e.g., connection **8**, as shown in FIG. 7) as the network route to the target datacenter (as shown in step **807** of FIG. 8). In the above embodiment, the mesh gateway may always use the direct connection to transmit the message to the target datacenter according to the user's preference even if the direct connection is either slow or down, and faster route is available.

(79) In an alternate embodiment, the mesh gateway (e.g., gateway **240**) may further determine that it does not have direct connection to the target datacenter (e.g., datacenter DC3 **250**, as shown in FIG. 7). In the above embodiment, upon determining that the direct connection to the target datacenter does not exist, the mesh gateway (e.g., gateway **240**) may follow the below steps to determine the network route for routing the protocol message to the target datacenter (e.g., datacenter DC3 **250**, as shown in FIG. 7). In step **808**, the mesh gateway (e.g., gateway **240**) determines one or more distance(s) to one or more next-hop datacenter(s) (e.g., DC1 **160**, as shown in FIG. 7).

(80) In the above embodiment, to determine relative distance to one or more next-hop datacenter(s) (e.g., DC1 **160**), the mesh gateway (e.g., gateway **240**) may use network coordinates to compute the relative distance between datacenters (datacenter DC2 **170** and datacenter DC1 **160**). Network coordinates provides an efficient way to measure network distance between these datacenter(s). In an example embodiment, control plane servers (e.g., server **166**) may configure routes other datacenter(s) based on priority and distance between datacenters. These configured routes are pushed to the gateway (e.g., gateway **230**) via xDS.

(81) In the above embodiment, as shown in step **809**, the gateway (e.g. gateway **240**, as shown in FIG. 7) may further determine one or more distance(s) to a target datacenter (e.g., DC3 **250**) as advertised by the one or more next-hop datacenter(s) (e.g., DC1 **160**). In the above embodiment, when the one or more next-hop datacenter(s) (e.g., DC1 **160**) advertises indirect path to target datacenter (e.g., DC3 **250**), the relative distance(s) to target datacenter (e.g., DC3 **250**) is included

in the payload.

(82) In the above embodiment, the mesh gateway (e.g., gateway **240**) may further calculate one or more total distance(s) (as shown in step **810**) to the target datacenter (e.g., DC3 **250**) based on the distance(s) to the one or more next-hop datacenter(s) (e.g., DC1 **160**) and the distance(s) to target datacenter (e.g., DC3 **250**) as advertised. In an example embodiment, the mesh gateway (e.g., gateway **240**) may calculate total distance(s) to the target datacenter (e.g., DC3 **250**) by adding the distance(s) to the one or more next-hop datacenter(s) and the distance(s) to target datacenter as advertised.

(83) In the above embodiment, under step **811**, the mesh gateway (e.g., gateway **240**) may further identify a network route with a shortest total distance (e.g., connections **13** and **9**, as shown in FIG. 7) to the target datacenter (e.g., DC3 **250**) from the one or more total distance(s) (e.g., connection route (**13**, **4**, and **5**) or connection route (**13** and **9**), as shown in FIG. 7) as calculated in step **810**. In an alternate embodiment, steps **808** to **811** are performed if upon determining that the routing configuration for the mesh gateway is a multi-hop routing, under step **804**. In the above embodiment, the mesh gateway may filter a reference to the first gateway device from the network route for routing the protocol message to prevent cyclical routes.

(84) In above embodiments, once the network route is determined, the mesh gateway (e.g., gateway **240**) may forward the protocol message to another gateway (e.g., gateway **210**) within the target datacenter (DC3 **250**) by routing the protocol message through the determined network route (e.g., connection route **13** and **9**, as shown in FIG. 7). Once the protocol message reaches the mesh gateway within the target datacenter, the mesh gateway may transmit the message to the target service internally within the target datacenter.

(85) In an example embodiment, a distance to target datacenter is calculated based on round trip times (RTTs) recorded between all datacenters using the WAN gossip protocol. In an example embodiment, for multi-hop network routes, RTTs for each hopping datacenter on a network route to the target datacenter is added to compute a total distance to the target datacenter. In the above embodiment, multiple distances are computed for multiple network routes to the target datacenter. In the above embodiment, a shortest distance to the target datacenter may be chosen as the network route. In certain embodiments, a topological distance model between all datacenters may be available to be used to compute and select the shortest route to the target datacenter. This model may be used to calculate distances and select a shortest route to the target datacenter where the shortest route also meets other criteria such as health, actually hosting an instance of the service of the target datacenter, or is on the path to another datacenter that hosts the instance of the service.

(86) In an example embodiment, the mesh gateway may use either a Local Area Network (LAN) or a Wide Area Network (WAN) Internet Protocol (IP) address to route the protocol message within the target datacenter. In an example embodiment, the mesh gateway may use at least one of a layer 4 and a layer 7 load balancing protocols for routing the protocol message.

(87) Numerous specific details are set forth herein to provide a thorough understanding of the claimed subject matter. However, those skilled in the art will understand that the claimed subject matter may be practiced without these specific details. In other instances, methods, apparatuses, or systems that would be known by one of ordinary skill have not been described in detail so as not to obscure claimed subject matter.

(88) Unless specifically stated otherwise, it is appreciated that throughout this specification discussions utilizing terms such as “processing,” “computing,” “calculating,” “determining,” and “identifying” or the like refer to actions or processes of a computing device, such as one or more computers or a similar electronic computing device or devices, that manipulate or transform data represented as physical electronic or magnetic quantities within memories, registers, or other information storage devices, transmission devices, or display devices of the computing platform.

(89) The system or systems discussed herein are not limited to any particular hardware architecture or configuration. A computing device can include any suitable arrangement of components that

provide a result conditioned on one or more inputs. Suitable computing devices include multi-purpose microprocessor-based computer systems accessing stored software that programs or configures the computing system from a general purpose computing apparatus to a specialized computing apparatus implementing one or more embodiments of the present subject matter. Any suitable programming, scripting, or other type of language or combinations of languages may be used to implement the teachings contained herein in software to be used in programming or configuring a computing device.

(90) Embodiments of the methods disclosed herein may be performed in the operation of such computing devices. The order of the blocks presented in the examples above can be varied—for example, blocks can be re-ordered, combined, and/or broken into sub-blocks. Certain blocks or processes can be performed in parallel.

(91) The use of “configured to” herein is meant as an open and inclusive language that does not foreclose devices adapted to or configured to perform additional tasks or steps. Where devices, systems, components or modules are described as being configured to perform certain operations or functions, such configuration can be accomplished, for example, by designing electronic circuits to perform the operation, by programming programmable electronic circuits (such as microprocessors) to perform the operation such as by executing computer instructions or code, or processors or cores programmed to execute code or instructions stored on a non-transitory memory medium, or any combination thereof. Processes can communicate using a variety of techniques including but not limited to conventional techniques for inter-process communications, and different pairs of processes may use different techniques, or the same pair of processes may use different techniques at different times.

(92) While the present subject matter has been described in detail with respect to specific embodiments thereof, it will be appreciated that those skilled in the art, upon attaining an understanding of the foregoing, may readily produce alterations to, variations of, and equivalents to such embodiments. Accordingly, it should be understood that the present disclosure has been presented for purposes of example rather than limitation, and does not preclude the inclusion of such modifications, variations, and/or additions to the present subject matter as would be readily apparent to one of ordinary skill in the art.

Claims

1. A method comprising: receiving, by a first gateway device, a protocol message directed to a service within a target datacenter from a plurality of datacenters; identifying, by the first gateway device, destination information of the service and the target datacenter; determining, by the first gateway device, a routing configuration for routing the protocol message to the target datacenter, wherein the routing configuration is a multi-hop routing; determining, by the first gateway device, a network route for routing the protocol message to the target datacenter based on the routing configuration; identifying, by the first gateway device, a route with a shortest total distance to the target datacenter, whereupon determining the routing configuration further comprises: determining, by the first gateway device, one or more distances to one or more next-hop datacenters and to the target datacenter as provided by the one or more next-hop datacenters, each distance being calculated based on the round trip time (RTT) recorded between all datacenters, wherein the RTT is calculated using a Wireless Area Network (WAN) gossip protocol; identifying, by the first gateway device, the network route with the shortest total distance to the target datacenter from the one or more total distances; and transmitting, by the first gateway device, the protocol message to a second gateway device associated with the target datacenter through the determined network route.
2. The method in accordance with claim 1, wherein the determining the one or more distances to the one or more next-hop datacenters further comprises calculating the one or more distances to the one or more next-hop datacenters using network coordinates of the datacenter associated with the

first gateway device and the one or more next-hop datacenters.

3. The method in accordance with claim 1, further comprising filtering, by the first gateway device, a reference to the first gateway device from the network route for routing the protocol message to prevent cyclical routes.

4. The method in accordance with claim 1, wherein the protocol message is a part of a Transport Layer Security (TLS)-Server Name Indication (SNI) handshake originated from a client service.

5. The method in accordance with claim 1, wherein the first gateway device identifies the destination information of the service and the target datacenter by sniffing into the protocol message.

6. The method in accordance with claim 1, wherein the destination information associated with the service and the target datacenter is included in a SNI header of the protocol message, and determined via TLS-SNI handshake.

7. The method in accordance with claim 1, wherein the first gateway device uses at least one of a Local Area Network (LAN) and a Wide Area Network (WAN) Internet Protocol (IP) address to route the protocol message to the service within the target datacenter.

8. The method in accordance with claim 1, wherein the first gateway device uses at least one of a layer 4 and a layer 7 load balancing protocols for routing the protocol message.
