

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388743
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Chanda; Anupam

Transaction controls for supplying control plane data to managed hardware forwarding element

Abstract

Some embodiments provide novel methods for controllers to communicate with managed hardware forwarding elements (MHFEs) in a transactional manner. The transactional communication methods of some embodiments ensure that an MHFE receives the entirety of a control plane update that a controller supplies to it, before the MHFE starts to modify its data plane forwarding data and operations. The transactional communication methods of some embodiments provide one or more transactional boundary controls to the controllers to define complete control plane data set updates. In some embodiments, the transactional controls ensure that an MHFE receives all of a control plane update before it starts to modify its data plane forwarding data. Controllers use one transactional control in some embodiments when they define logical forwarding elements (e.g., logical switches or routers) on the MHFEs.

Inventors:	Chanda; Anupam (San Jose, CA)
Applicant:	VMware LLC (Palo Alto, CA)
Family ID:	59020294
Assignee:	VMware LLC (Palo Alto, CA)
Appl. No.:	15/959261
Filed:	January 05, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20180241672 A1	Aug. 23, 2018

Related U.S. Application Data

continuation parent-doc US 14970463 20151215 US 9998375 child-doc US 15959261

Publication Classification

Int. Cl.: H04L12/00 (20060101); H04L45/42 (20220101); H04L45/745 (20220101); H04L49/00 (20220101)

U.S. Cl.:

CPC H04L45/42 (20130101); H04L45/745 (20130101); H04L49/70 (20130101);

Field of Classification Search

CPC: H04L (45/42); H04L (45/745); H04L (49/70)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5526484	12/1995	Casper et al.	N/A	N/A
5751967	12/1997	Raab et al.	N/A	N/A
6243394	12/2000	Deng	N/A	N/A
6640251	12/2002	Wiget et al.	N/A	N/A
6850529	12/2004	Wong et al.	N/A	N/A
7463639	12/2007	Rekhter	N/A	N/A
7933198	12/2010	Pan	N/A	N/A
7937438	12/2010	Miller et al.	N/A	N/A
8046456	12/2010	Miller et al.	N/A	N/A
8161095	12/2011	Manion et al.	N/A	N/A
8345688	12/2012	Zhou et al.	N/A	N/A
8386642	12/2012	Elzur	N/A	N/A
8589919	12/2012	Smith et al.	N/A	N/A
8874876	12/2013	Bhadra et al.	N/A	N/A
8897134	12/2013	Kern et al.	N/A	N/A
8943490	12/2014	Jain et al.	N/A	N/A
8964528	12/2014	Casado et al.	N/A	N/A
9014181	12/2014	Lakshman et al.	N/A	N/A
9032095	12/2014	Traina et al.	N/A	N/A
9100285	12/2014	Choudhury et al.	N/A	N/A
9130870	12/2014	Swierk et al.	N/A	N/A
9154433	12/2014	Koponen et al.	N/A	N/A
9178833	12/2014	Koponen et al.	N/A	N/A
9306843	12/2015	Koponen et al.	N/A	N/A
9319375	12/2015	Gross, IV et al.	N/A	N/A
9331940	12/2015	Balus et al.	N/A	N/A
9369426	12/2015	Koponen et al.	N/A	N/A
9397946	12/2015	Yadav	N/A	N/A
9455901	12/2015	Davie et al.	N/A	N/A
9485149	12/2015	Traina et al.	N/A	N/A
9577927	12/2016	Hira et al.	N/A	N/A
9621461	12/2016	Sun	N/A	N/A

9633040	12/2016	Lee	N/A	N/A
9667541	12/2016	Song	N/A	N/A
9699070	12/2016	Davie et al.	N/A	N/A
9755965	12/2016	Yadav et al.	N/A	N/A
9819581	12/2016	Chanda et al.	N/A	N/A
9847938	12/2016	Chanda et al.	N/A	N/A
9917799	12/2017	Chanda	N/A	N/A
9923815	12/2017	Assarpour et al.	N/A	N/A
9942058	12/2017	Chanda et al.	N/A	N/A
9948577	12/2017	Chanda	N/A	N/A
9967182	12/2017	Chanda et al.	N/A	N/A
9979593	12/2017	Chanda et al.	N/A	N/A
9992112	12/2017	Chanda	N/A	N/A
9998324	12/2017	Chanda et al.	N/A	N/A
9998375	12/2017	Chanda	N/A	N/A
10153965	12/2017	Davie et al.	N/A	N/A
10182035	12/2018	Basler	N/A	N/A
10200343	12/2018	Chanda et al.	N/A	N/A
10230576	12/2018	Chanda et al.	N/A	N/A
10250553	12/2018	Chanda et al.	N/A	N/A
10263828	12/2018	Chanda et al.	N/A	N/A
10313186	12/2018	Wang et al.	N/A	N/A
10411912	12/2018	Chanda et al.	N/A	N/A
10447618	12/2018	Chanda	N/A	N/A
10554484	12/2019	Chanda et al.	N/A	N/A
10659431	12/2019	Basler	N/A	N/A
10764111	12/2019	Chanda et al.	N/A	N/A
10805152	12/2019	Chanda et al.	N/A	N/A
10924386	12/2020	Davie et al.	N/A	N/A
11005683	12/2020	Chanda et al.	N/A	N/A
11032234	12/2020	Chanda et al.	N/A	N/A
11095513	12/2020	Wang et al.	N/A	N/A
11196682	12/2020	Chanda	N/A	N/A
11245621	12/2021	Chanda et al.	N/A	N/A
11368431	12/2021	Basler	N/A	N/A
11502898	12/2021	Chanda et al.	N/A	N/A
11522788	12/2021	Davie et al.	N/A	N/A
2002/0001310	12/2001	Mai et al.	N/A	N/A
2002/0065919	12/2001	Taylor et al.	N/A	N/A
2002/0093952	12/2001	Gonda	N/A	N/A
2003/0041170	12/2002	Suzuki	N/A	N/A
2003/0163645	12/2002	Tremblay et al.	N/A	N/A
2004/0037279	12/2003	Zelig et al.	N/A	N/A
2004/0052216	12/2003	Roh	N/A	N/A
2004/0267866	12/2003	Carollo et al.	N/A	N/A
2005/0163115	12/2004	Dontu et al.	N/A	N/A
2005/0262132	12/2004	Morita et al.	N/A	N/A
2006/0092940	12/2005	Ansari et al.	N/A	N/A
2007/0028039	12/2006	Gupta et al.	N/A	N/A
2007/0115961	12/2006	Dorenbosch et al.	N/A	N/A

2008/0189769	12/2007	Casado et al.	N/A	N/A
2008/0215586	12/2007	Pruet	N/A	N/A
2009/0135823	12/2008	Liu et al.	N/A	N/A
2009/0006603	12/2008	Duponchel et al.	N/A	N/A
2009/0150527	12/2008	Tripathi et al.	N/A	N/A
2009/0161547	12/2008	Riddle et al.	N/A	N/A
2011/0026521	12/2010	Gamage et al.	N/A	N/A
2011/0075674	12/2010	Li et al.	N/A	N/A
2011/0090911	12/2010	Hao et al.	N/A	N/A
2011/0131361	12/2010	Takura	N/A	N/A
2011/0158086	12/2010	Sakauchi et al.	N/A	N/A
2011/0206047	12/2010	Donthamsetty et al.	N/A	N/A
2011/0286326	12/2010	Awano	N/A	N/A
2011/0299537	12/2010	Saraiya et al.	N/A	N/A
2011/0317559	12/2010	Kern et al.	N/A	N/A
2012/0011106	12/2011	Reid et al.	N/A	N/A
2012/0084406	12/2011	Kumbalimutt	709/220	H04L 41/5054
2012/0147898	12/2011	Koponen et al.	N/A	N/A
2012/0215876	12/2011	Ohtake	N/A	N/A
2012/0236761	12/2011	Yang et al.	N/A	N/A
2012/0278802	12/2011	Nilakantan et al.	N/A	N/A
2012/0290694	12/2011	Marl et al.	N/A	N/A
2012/0303835	12/2011	Kempf et al.	N/A	N/A
2013/0044636	12/2012	Koponen et al.	N/A	N/A
2013/0044641	12/2012	Koponen et al.	N/A	N/A
2013/0051399	12/2012	Zhang et al.	N/A	N/A
2013/0054773	12/2012	Onoue	N/A	N/A
2013/0058208	12/2012	Pfaff et al.	N/A	N/A
2013/0058225	12/2012	Casado	370/250	H04L 12/4633
2013/0058229	12/2012	Casado et al.	N/A	N/A
2013/0058350	12/2012	Fulton	N/A	N/A
2013/0058351	12/2012	Casado et al.	N/A	N/A
2013/0060929	12/2012	Koponen et al.	N/A	N/A
2013/0103817	12/2012	Koponen et al.	N/A	N/A
2013/0103818	12/2012	Koponen et al.	N/A	N/A
2013/0114466	12/2012	Koponen et al.	N/A	N/A
2013/0121209	12/2012	Padmanabhan	370/255	H04L 67/1008
2013/0132533	12/2012	Padmanabhan	709/220	H04L 45/74
2013/0223454	12/2012	Dunbar et al.	N/A	N/A
2013/0287026	12/2012	Davie	N/A	N/A
2013/0315246	12/2012	Zhang et al.	N/A	N/A
2013/0322453	12/2012	Allan	N/A	N/A
2013/0329548	12/2012	Nakil et al.	N/A	N/A
2013/0336134	12/2012	Bao	370/252	H04L 41/046
2014/0029451	12/2013	Nguyen	N/A	N/A
2014/0029618	12/2013	Janardhanan	N/A	N/A

2014/0071986	12/2013	Isobe	N/A	N/A
2014/0101467	12/2013	Jubran et al.	N/A	N/A
2014/0195666	12/2013	Dumitriu	709/223	H04L 12/4625
2014/0201738	12/2013	Choi et al.	N/A	N/A
2014/0229605	12/2013	Besser	N/A	N/A
2014/0247753	12/2013	Koponen	370/255	H04L 45/72
2014/0269290	12/2013	Cors et al.	N/A	N/A
2014/0269683	12/2013	Bhagavathiperumal et al.	N/A	N/A
2014/0269709	12/2013	Benny et al.	N/A	N/A
2014/0301391	12/2013	Krishnan et al.	N/A	N/A
2014/0348161	12/2013	Koponen	370/389	H04L 41/0654
2014/0362775	12/2013	Steiner	370/329	G06F 9/45558
2015/0009992	12/2014	Zhang	N/A	N/A
2015/0016287	12/2014	Ganichev et al.	N/A	N/A
2015/0016469	12/2014	Ganichev et al.	N/A	N/A
2015/0100560	12/2014	Davie et al.	N/A	N/A
2015/0100675	12/2014	Davie et al.	N/A	N/A
2015/0100704	12/2014	Davie et al.	N/A	N/A
2015/0103838	12/2014	Zhang et al.	N/A	N/A
2015/0124586	12/2014	Pani	N/A	N/A
2015/0124809	12/2014	Edsall	370/390	H04L 47/20
2015/0124821	12/2014	Chu et al.	N/A	N/A
2015/0195178	12/2014	Bhattacharya et al.	N/A	N/A
2015/0215189	12/2014	Lim	N/A	N/A
2015/0281075	12/2014	Park et al.	N/A	N/A
2015/0312117	12/2014	Behera	370/241	H04L 43/04
2015/0326425	12/2014	Natarajan et al.	N/A	N/A
2015/0326524	12/2014	Tankala et al.	N/A	N/A
2015/0334696	12/2014	Gu	718/1	H04L 61/2015
2015/0372906	12/2014	Tirat	N/A	N/A
2015/0379150	12/2014	Duda	707/736	G06F 16/955
2015/0381428	12/2014	Ong	N/A	N/A
2016/0014039	12/2015	Reddy et al.	N/A	N/A
2016/0014073	12/2015	Reddy et al.	N/A	N/A
2016/0173535	12/2015	Barabash	726/1	H04L 41/0893
2016/0197824	12/2015	Lin et al.	N/A	N/A
2016/0212222	12/2015	Bultema et al.	N/A	N/A
2016/0232019	12/2015	Shah et al.	N/A	N/A
2016/0308690	12/2015	Chanda et al.	N/A	N/A
2016/0337896	12/2015	Rajagopalan	N/A	H04L 41/0893
2016/0352613	12/2015	Narayanan	N/A	H04L 45/021

2016/0352633	12/2015	Kapadia et al.	N/A	N/A
2016/0380812	12/2015	Chanda et al.	N/A	N/A
2017/0034002	12/2016	Sinn	N/A	N/A
2017/0034051	12/2016	Chanda et al.	N/A	N/A
2017/0034052	12/2016	Chanda et al.	N/A	N/A
2017/0034053	12/2016	Chanda et al.	N/A	N/A
2017/0063608	12/2016	Wang et al.	N/A	N/A
2017/0085502	12/2016	Biruduraju	N/A	N/A
2017/0093617	12/2016	Chanda et al.	N/A	N/A
2017/0093618	12/2016	Chanda et al.	N/A	N/A
2017/0093636	12/2016	Chanda et al.	N/A	N/A
2017/0093646	12/2016	Chanda et al.	N/A	N/A
2017/0093686	12/2016	Uttaro et al.	N/A	N/A
2017/0093758	12/2016	Chanda	N/A	N/A
2017/0126615	12/2016	Chanda et al.	N/A	N/A
2017/0171055	12/2016	Wang et al.	N/A	N/A
2017/0171077	12/2016	Chanda	N/A	N/A
2017/0171078	12/2016	Chanda	N/A	N/A
2017/0171113	12/2016	Chanda	N/A	N/A
2017/0208097	12/2016	Kirby et al.	N/A	N/A
2017/0272316	12/2016	Johnson et al.	N/A	N/A
2017/0317928	12/2016	Gude et al.	N/A	N/A
2017/0339054	12/2016	Yadav	N/A	H04L 41/5054
2017/0366446	12/2016	Davie et al.	N/A	N/A
2018/0007004	12/2017	Basler	N/A	N/A
2018/0007005	12/2017	Chanda et al.	N/A	N/A
2018/0013687	12/2017	Jin	N/A	H04L 47/36
2018/0026895	12/2017	Wang et al.	N/A	N/A
2018/0183730	12/2017	Chanda	N/A	N/A
2018/0219699	12/2017	Chanda et al.	N/A	N/A
2018/0241622	12/2017	Chanda et al.	N/A	N/A
2018/0248796	12/2017	Chanda et al.	N/A	N/A
2019/0089622	12/2018	Davie et al.	N/A	N/A
2019/0141011	12/2018	Basler	N/A	N/A
2019/0207904	12/2018	Chanda et al.	N/A	N/A
2019/0238393	12/2018	Chanda et al.	N/A	N/A
2019/0260637	12/2018	Wang et al.	N/A	N/A
2019/0356512	12/2018	Chanda et al.	N/A	N/A
2020/0021541	12/2019	Chanda	N/A	N/A
2020/0366553	12/2019	Chanda et al.	N/A	N/A
2020/0396206	12/2019	Basler	N/A	N/A
2021/0243112	12/2020	Davie et al.	N/A	N/A
2022/0078111	12/2021	Chanda et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1154601	12/2000	EP	N/A
1653688	12/2005	EP	N/A

2003069609	12/2002	JP	N/A
2005112390	12/2004	WO	N/A
2008095010	12/2007	WO	N/A

OTHER PUBLICATIONS

Al-Fares, Mohammad, et al., "A Scalable, Commodity Data Center Network Architecture," SIGCOMM '08, Aug. 17-22, 2008, 12 pages, Seattle, Washington, USA. cited by applicant

Graubner, Pablo, et al., "Cloud Computing: Energy-Efficient Virtual Machine Consolidation," IT Professional, Mar. 2013, 7 pages, vol. 15, Issue 2, IEEE. cited by applicant

Pfaff, Ben, et al., "The Open vSwitch Database Management Protocol," draft-pfaff-ovsdb-proto-00, Aug. 20, 2012, 34 pages, available at <https://tools.ietf.org/html/draft-pfaff-ovsdb-proto-00>. cited by applicant

Primary Examiner: Beharry; Noel R

Assistant Examiner: Preval; Lionel

Attorney, Agent or Firm: Quarles & Brady LLP

Background/Summary

CLAIM OF BENEFIT TO PRIOR APPLICATIONS (1) This application is a continuation application of U.S. patent application Ser. No. 14/970,463, filed Dec. 15, 2015, and now published as U.S. Patent Publication 2017/0171077. U.S. patent application Ser. No. 14/970,463, now published as U.S. Patent Publication 2017/0171077 is incorporated herein by reference.

BACKGROUND

(1) Today, several network control systems configure software forwarding elements executing on host computers to create logical forwarding elements that span multiple host computers. Such logical forwarding elements allow multiple isolated logical networks to be created on a shared physical network infrastructure. In recent years, some network control systems have been adapted to configure hardware forwarding elements, such as top of the rack (TOR) switches. In many such cases, the TORs have configuration databases that are defined according to the hardware VTEP (VXLAN Tunnel Endpoint) schema. Network controllers of such network control systems communicate with such TOR configuration databases using the OVSdb protocol. The OVSdb protocol and hardware VTEP schema are defined by the Open vSwitch organization.

(2) The hardware VTEP schema defines various tables to exchange control plane information between a controller and a TOR. However, neither this schema nor the OVSdb protocol require the controllers to use transactional controls in communicating with the TORs. As such, on receiving a portion of a configuration from a controller, the TOR might start modifying its data plane operation. This may cause problems as the received configuration data portion might provide a view of the logical network that is incomplete and in some cases, inconsistent with either an earlier logical network view upon which the TOR previously based its data plane operation, or with the current logical network view that the controller is trying to push to the TOR. Therefore, there is a need in the art for transactional controls for the network controllers to use to allow TORs to update their data plane operations based on complete control plane views of the logical networks.

BRIEF SUMMARY

(3) Some embodiments provide novel methods for controllers to communicate with managed hardware forwarding elements (MHFEs) in a transactional manner. The transactional

communication methods of some embodiments ensure that an MHFE receives the entirety of a control plane update that a controller supplies to it, before the MHFE starts to modify its data plane forwarding data and operations.

(4) The network controllers in some embodiments manage the MHFEs to create one or more logical networks that span shared physical forwarding elements, including the MHFEs. In some embodiments, the shared physical forwarding elements also include software forwarding elements executing on host computers, on which multiple compute nodes (e.g., virtual machines, containers, etc.) execute. The transactional communication methods of some embodiments ensure that the MHFEs update their data plane operations based on complete control plane views of logical networks.

(5) The transactional communication methods of some embodiments provide one or more transactional boundary controls to the controllers to define complete control plane data set updates. In some embodiments, the transactional controls ensure that an MHFE receives all of a control plane update before it starts to modify its data plane forwarding data. Controllers use one transactional control in some embodiments when they define logical forwarding elements (e.g., logical switches or routers) for the first time on the MHFEs.

(6) In some embodiments, one configuration data tuple that is needed to create a logical forwarding element on an MHFE is a data tuple that binds (i.e., associates) the logical forwarding element with a physical port of the MHFE (e.g., a port of a hardware top-of-rack switch). This configuration data tuple is referred to below as the LAN (local area network) binding tuple. In pushing configuration data to an MHFE to define a logical forwarding element (LFE) on the MHFE, the controller in some embodiments supplies other configuration data tuples to the MHFE before the LAN binding tuple. In some of these embodiments, the controller supplies the LAN binding tuple as the last configuration data tuple to the MHFE. This is because the MHFE in some embodiments starts to create its data plane forwarding records for the LFE once the MHFE receives the LAN binding tuple.

(7) In some embodiments, the MHFEs use the hardware-VTEP schema. In this schema, there is a Logical_Switch table where the controller creates logical switches, and a Ucast_Macs_Remote table, where the controller provides the forwarding information for those logical switches. Another table in this schema is the Physical_Port table, which includes a vlan_bindings column that binds specific VLANs on the specific physical ports of the hardware VTEP device (i.e., of an MHFE) to corresponding logical switches. To achieve transactional semantics for a logical switch, the controllers of some embodiments propagate all modifications to the Logical_Switch table and the Ucast_Macs_Remote table before updating the vlan_bindings column in the Physical_Port table for a new logical switch that the controller is defining on the MHFE.

(8) When updating a logical network's control plane data, some embodiments (1) delete the vlan_binding tuple in the Physical_Port table to unbind the logical switch from the physical network, (2) update the configuration data in one or more MHFE's tables (e.g., forwarding information in Ucast_Macs_Remote table), and then (3) recreate the binding in the vlan_bindings column. This approach requires the logical switch to stop forwarding during the time between when the vlan_binding tuple is deleted and when it is added back. In other words, during this time, there will be a data plane outage.

(9) Accordingly, other embodiments use other transaction controls to update the control plane configuration data of a logical switch that is already defined on an MHFE, instead of unbinding the logical switch from the MHFE. In some embodiments, an MHFE creates a lock for a logical forwarding element that is defined on it. In some of these embodiments, the controller can “steal” this lock (e.g., can take this lock even when another MHFE module is using it to access the hardware VTEP database) during a period in which the controller updates the configuration data of a logical forwarding element (LFE). While the controller has stolen the LFE's lock, no MHFE module examines its configuration data storage to detect changes to the configuration data and to

update its data plane forwarding records based on such detected changes. Once the controller completes its control plane configuration update for a LFE implemented by the MHFE, the controller returns the lock for the LFE (i.e., relinquishes the lock), at which point an MHFE module can request the lock for the LFE, can detect changes to the LFE's control plane configuration, and then can update its data plane forwarding data based on these detected changes.

(10) In the embodiments where the MHFE uses the hardware VTEP schema, this schema supports creation of a “lock” for one or more records in the hardware VTEP database. Once a lock is created for a record, a hardware VTEP database client can use the following commands to communicate with a hardware VTEP database server (for the hardware VTEP database) vis-à-vis the lock: lock request, lock steal request, and lock release. Lock request is a request to obtain the lock for a data record in order to read and/or write to the data record. When another database client has a lock for the data record, a lock request does not cause the database server to remove the lock from the other database client to provide the lock to the database client that requested the lock. A lock steal request, however, does cause the database server to remove the lock from another database client that currently has the lock in order to provide the lock to the database client that provided the lock steal request. When a database client is done with a lock, it provides a lock release command to the database server to release the lock.

(11) Some embodiments create a lock for a logical switch record in the Logical_Switch table. Once a binding in the vlan_bindings column is established, a database client on the MHFE acquires a lock for the logical switch by providing the logical switch's identifier along with the lock request to the MHFE's hardware VTEP database server. Subsequently, when the controller wishes to change the logical switch's configuration (e.g., the logical switch's control plane data records), the controller's hardware VTEP database client provides a lock steal request and the logical switch's identifier to the database server.

(12) The MHFE's database server provides a notification to its database client that the lock has been stolen, which indicates a forthcoming change to the forwarding information. The MHFE's database client tries to re-acquire the lock. However, its attempts fail as long as the controller is holding the stolen lock while it makes changes. The forwarding plane continues to work according to the forwarding information of the last transaction, i.e., there is no outage. After the controller is done with making its changes, it releases the lock for the logical switch. At this point, the MHFE's database client's lock re-acquisition operation succeeds. When the MHFE database client is able to re-acquire the lock, it applies all the changes that the controller made to the forwarding table to its forwarding plane. Thus, in this manner, the forwarding plane for the logical switch is transactionally updated in some embodiments.

(13) Instead of using the lock-based approach, other embodiments define a configuration data tuple that the controller uses to explicitly inform the MHFE that the control plane configuration data is being updated and hence might not provide a consistent view of a logical forwarding element or a logical network. For instance, instead of using the lock-based approach, some embodiments change the hardware VTEP schema to add a Boolean column to the Logical_Switch table to indicate whether the logical switch's information in the hardware VTEP database is consistent from the controller's point of view. This column is referred to below as the state_consistent column.

(14) In some embodiments, the controller initializes this column to False while creating a logical switch. The controller is free to update the vlan_binding column at any point for this logical switch. After supplying all the data tuples for the logical switch (e.g., all the data tuples to the Ucast_Macs_Remote table, Mcast_Macs_Remote table, and Physical_Port table), the controller updates the state_consistent column to True. Once that happens, the MHFE database client modifies the MHFE's forwarding tables in the data plane based on the supplied logical switch configuration tuples.

(15) When the controller wants to modify the MHFE's control plane configuration for a logical switch, the controller first changes the state_consistent column to False. The MHFE agent monitors

this column. When the column becomes False, the MHFE's forwarding records in the data plane continue to work in its last transactional state. The controller makes the changes to the MHFE's hardware VTEP database, and then changes the state_consistent column to True. At this point, the MHFE's database client applies all changes to its forwarding records in the data planes.

(16) The preceding Summary is intended to serve as a brief introduction to some embodiments of the invention. It is not meant to be an introduction or overview of all inventive subject matter disclosed in this document. The Detailed Description that follows and the Drawings that are referred to in the Detailed Description will further describe the embodiments described in the Summary as well as other embodiments. Accordingly, to understand all the embodiments described by this document, a full review of the Summary, Detailed Description, the Drawings and the Claims is needed. Moreover, the claimed subject matters are not to be limited by the illustrative details in the Summary, Detailed Description and the Drawing.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The novel features of the invention are set forth in the appended claims. However, for purposes of explanation, several embodiments of the invention are set forth in the following figures.
- (2) FIG. 1 illustrates four of tables that are defined in the OVSdb schema.
- (3) FIG. 2 illustrates a modified Logical Switch table with a state_consistent column.
- (4) FIG. 3 conceptually illustrates a network control system that implements the MHFE management methods of some embodiments of the invention.
- (5) FIG. 4 conceptually illustrates how a network controller communicates with a software switch and a hardware switch in some embodiments.
- (6) FIG. 5 conceptually illustrates a process that a physical controller performs in some embodiments to supply the definition of a logical forwarding element to an MHFE in a transactional manner.
- (7) FIG. 6 conceptually illustrates a process that a TOR's OVSdb server performs to define a logical switch on the TOR.
- (8) FIG. 7 conceptually illustrate the processes that the TOR agent, OVSdb server and the OVSdb client perform in some embodiments during and after a period in which the TOR agent modifies the records of a logical switch in the OVS database.
- (9) FIG. 8 conceptually illustrates a process that a TOR agent performs in some embodiments to supply define a logical switch or modify a logical switch definition in a transactional manner for a TOR.
- (10) FIG. 9 conceptually illustrates a process that a hardware switch's OVSdb server performs to create a logical switch on the hardware switch, or to update the configuration of a previously created logical switch on the hardware switch.
- (11) FIG. 10 conceptually illustrates a computer system that implements processes of some embodiments of the invention.

DETAILED DESCRIPTION

- (12) In the following detailed description of the invention, numerous details, examples, and embodiments of the invention are set forth and described. However, it will be clear and apparent to one skilled in the art that the invention is not limited to the embodiments set forth and that the invention may be practiced without some of the specific details and examples discussed.
- (13) Some embodiments provide novel methods for controllers to communicate with managed hardware forwarding elements (MHFEs) in a transactional manner. The transactional communication methods of some embodiments ensure that an MHFE receives the entirety of a control plane update that a controller supplies to it, before the MHFE starts to modify its data plane

forwarding data and operations. Examples of MHFEs include top-of-rack (TOR) switches.

(14) The network controllers in some embodiments manage the MHFEs to create one or more logical forwarding elements that span shared physical forwarding elements, including the MHFEs. In some embodiments, the shared physical forwarding elements also include software forwarding elements executing on host computers, on which multiple data compute nodes (DCNs, such as virtual machines, containers, etc.) execute. The logical forwarding elements allow multiple logical networks to be created on a shared physical compute and network infrastructure.

(15) Each logical network's set of one or more logical forwarding elements (LFEs) can connect DCNs (e.g., virtual machines, containers, etc.) that execute on different host computers. Each LFE is an abstract construct that conceptually spans multiple forwarding elements (e.g., software forwarding elements, SFEs) to connect DCNs on multiple different hosts to each other. In some embodiments, overlay tunnel connections between the hosts facilitate the creation of the LFEs. Each logical network can isolate the traffic exchanged between its DCNs from the traffic from DCNs of other logical networks. Many examples of logical networks and logical forwarding elements are described in U.S. Published Patent Application 2013/0058335, U.S. Pat. No. 8,958,298, U.S. Published Patent Application 2015/0063360, and U.S. Published Patent Application 2015/0263946.

(16) As used in this document, data messages or packets generically refer to a collection of bits in a particular format sent across a network. One of ordinary skill in the art will recognize that the term data message or packet may be used herein to refer to various formatted collections of bits that may be sent across a network, such as Ethernet frames, IP packets, TCP segments, UDP datagrams, etc. Also, as used in this document, references to L2, L3, L4, and L7 layers (or layer 2, layer 3, layer 4, layer 7) are references respectively to the second data link layer, the third network layer, the fourth transport layer, and the seventh application layer of the OSI (Open System Interconnection) layer model.

(17) The transactional communication methods of some embodiments ensure that the MHFEs update their data plane operations based on complete control plane views of logical networks. The transactional communication methods of some embodiments provide one or more transactional boundary controls to the controllers to define complete control plane data set updates. In some embodiments, the transactional controls ensure that an MHFE receives all of a control plane update before it starts to modify its data plane forwarding data.

(18) Controllers use one transactional control in some embodiments when they define logical forwarding elements (e.g., logical switches or routers) for the first time on the MHFEs. In some embodiments, one configuration data tuple that is needed to create a logical forwarding element on an MHFE is a data tuple that binds (i.e., associates) the logical forwarding element with a physical port of the MHFE (e.g., a port of a TOR switch). This configuration data tuple is referred to below as the LAN (local area network) binding tuple. In pushing configuration data to an MHFE to define a logical forwarding element (LFE) on the MHFE, the controller in some embodiments supplies other configuration data tuples to the MHFE before the LAN binding tuple. In some of these embodiments, the controller supplies the LAN binding tuple as the last configuration data tuple to the MHFE. This is because the MHFE in some embodiments starts to create its data plane forwarding records for the LFE once the MHFE receives the LAN binding tuple.

(19) The MHFEs in some embodiments contain configuration databases and execute database servers with which the network controllers can interact through a database protocol. For instance, in some embodiments, the MHFE configuration database is a hardware VTEP database and the communication protocol is an OVSdb (an Open Virtual Switch (OVS) database) protocol. In some embodiments, the hardware VTEP schema for the MHFE specifies multiple MHFE database tables to which a controller can write. In the hardware VTEP schema, the LAN binding tuple is defined in one column of the Physical_Port table of this schema.

(20) FIG. 1 illustrates this table along with three other tables in the hardware VTEP schema. The

illustrated tables are the Ucast_Macs_Remote table **105**, Mcast_Macs_Remote table **110**, the Logical_Switch table **115**, and the Physical_Port table **120**. In an MHFE's database, the Logical_Switch table **115** contains a record for each logical switch implemented by the MHFE. For each logical switch that the MHFE implements, a controller in some embodiments creates a record in the Logical_Switch table **115**. As shown, each logical switch's record includes (1) a UUID (universally unique identifier) **122** tuple that provides a universal unique identifier for the logical switch, (2) a description **124** that provides a textual description of the logical switch, (3) a name tuple **126** that provides the name of the logical switch, and (4) a tunnel_key **128** that provides the VNI (VXLAN network identifier) of the logical switch.

(21) The Ucast_Macs_Remote table **105** provides the forwarding information for the logical switches. The Ucast table provides the unicast MAC (media access control) addresses for the DCNs (VMs, containers, etc.) associated with each logical switch. Optionally, this table also provides the IP (Internet Protocol) address of each of these VMs. This table also provides the VTEP IPs for each logical switch.

(22) As shown, each record in the Ucast_Macs_Remote table **105** includes (1) a logical switch field **132** that specifies the logical switch with which a port of an end machine is associated (in the case that the table is instantiated on an L3 MHFE, this field also specifies the logical router ports associated with the logical switch), (2) a MAC address field **134** that specifies the corresponding MAC address of the port (the MAC address of the end machine's port associated with the logical switch's port), (3) an optional field **136** that can include an IP address associated with the MAC address, and (4) a locator field **138** that specifies the IP address of the VTEP for the corresponding MAC address.

(23) Because of the locator field **138**, the Ucast table **105** is referred to as a tunnel endpoint locator table or a tunnel endpoint table. For an L2 MHFE, the controllers in some embodiments supply the Ucast table with the physical locator addresses (i.e., IP addresses) of the MFEs (hardware and software) that implement (1) the different logical switches' ports that are associated with the end machines of the logical network, and (2) the logical router ports that receive the L3 packets from the MHFE. The Ucast table **105** specifies the next destination of a unicast packet with a unique destination MAC address. By locating the endpoints, the L2 MHFE is able to establish tunnels between the MHFE and other MFEs and exchange the network data through the established tunnels.

(24) The Mcast_Macs_Remote table **110** is another tunnel endpoint locator table. The Mcast table **110** specifies the next destination of a broadcast, multicast, or unknown unicast packet that does not have a unique destination MAC address. As shown, the table **110** has three different fields **142**, **144**, and **146** that are similar to the fields **132**, **134**, and **138** of Ucast table **105**, with the exception that Mcast table **110** is for the destination MAC addresses that are not known to the MHFE or the destination MAC addresses of multicast and broadcast packets in some embodiments. The controllers in some embodiments configure the locator column **146** of the Mcast tables **110** with IP addresses of service nodes that process BUM (broadcast, unknown, and multicast) packets. This way, the MHFEs forward any BUM packet to service nodes for processing.

(25) The Physical_Port table **120** is a table that is specified by the MHFEs and read by the controllers, except that the controller updates the VLAN_bindings column **158** to bridge VLANs (virtual LANs) on specific physical ports of the MHFE to logical switches. As shown, this table **120** includes three columns **152**, **154** and **156** that define the port UUID, the port description, and the port name. The vlan_bindings column **158** is a map of a VLAN to a logical switch UUID. This column **158** establishes the logical switch to physical port association (for given VLANs).

(26) Updating the VLAN binding column is the transaction that causes a logical switch to bridge to the physical network. Some embodiments control this update in order to provide a transactional boundary in the pushed transaction data. Specifically, to achieve transactional semantics for a logical switch, the controllers of some embodiments propagate all configuration data to the MHFE

(e.g., all the data to the Logical_Switch table **115**, the Ucast_Macs_Remote table **105**, and the Mcast_Macs_Remote table **110**) before updating the vlan_bindings column **158** in the Physical_Port table **120**.

(27) When updating a logical network's control plane data, some embodiments (1) delete the vlan_binding tuple in the Physical_Port table **120** to unbind the logical switch from the physical network, (2) update the configuration data in one or more MHFE's tables (e.g., forwarding information in Ucast_Macs_Remote table **105**), and then (3) recreate the binding in the vlan_bindings column. This approach requires the logical switch to stop forwarding during the time between when the vlan_binding tuple is deleted and when it is added back. In other words, during this time, there will be a data plane outage.

(28) Accordingly, other embodiments use other transaction controls to update the control plane configuration data of a logical switch that is already defined on an MHFE, instead of unbinding the logical switch from the MHFE. In some embodiments, an MHFE creates a lock for a logical forwarding element that is defined on it. In some of these embodiments, the controller can “steal” this lock (e.g., can take this lock even when another MHFE module is using it to access the hardware VTEP database) during a period in which the controller updates the configuration data of a logical forwarding element (LFE). While the controller has stolen the LFE's lock, no MHFE module examines its configuration data storage to detect changes to the configuration data and to update its data plane forwarding records based on such detected changes. Once the controller completes its control plane configuration update for a LFE implemented by the MHFE, the controller returns the lock for the LFE (i.e., relinquishes the lock), at which point an MHFE module can request the lock for the LFE, can detect changes to the LFE's control plane configuration, and then can update its data plane forwarding data based on these detected changes.

(29) In the embodiments where the MHFE uses the hardware VTEP schema, this schema supports creation of a “lock” for one or more records in the hardware VTEP database. Once a binding in the vlan_bindings column is established, a database server executing on the MHFE creates a lock for the logical switch in some embodiments. The database server creates this lock at the direction of a database client on the controller in some embodiments. In some embodiments, this lock has an identifier that identifies the logical switch.

(30) Once a lock is created for a record, a hardware VTEP database client can use the following commands to communicate with a hardware VTEP database server (for the hardware VTEP database) vis-à-vis the lock: lock request, lock steal request, and lock release. Lock request is a request to obtain the lock for a data record in order to read and/or write to the data record. When another database client has a lock for the data record, a lock request does not cause the database server to remove the lock from the other database client to provide the lock to the database client that requested the lock. A lock steal request, however, does cause the database server to remove the lock from another database client that currently has the lock in order to provide the lock to the database client that provided the lock steal request. When a database client is done with a lock, it provides a lock release command to the database server to release the lock.

(31) Some embodiments create a lock for a logical switch record in the Logical_Switch table. Once the binding in the vlan_bindings column is established and the lock is created, a database client on the MHFE acquires a lock for the logical switch by providing the logical switch's identifier along with the lock request to the MHFE's hardware VTEP database server. Subsequently, when the controller wishes to change the logical switch's configuration (e.g., the logical switch's control plane data records), the controller's hardware VTEP database client provides a lock steal request and the logical switch's identifier to the database server.

(32) The MHFE's database server provides a notification to its database client that the lock has been stolen, which indicates a forthcoming change to the forwarding information. The MHFE's database client tries to re-acquire the lock. However, its attempts fail as long as the controller is holding the stolen lock while it makes changes. The forwarding plane continues to work according

to the forwarding information of the last transaction, i.e., there is no outage. After the controller is done with making its changes, it releases the lock for the logical switch. At this point, the MHFE's database client's lock re-acquisition operation succeeds. When the MHFE database client is able to re-acquire the lock, it applies all the changes that the controller made to the forwarding table to its forwarding plane. Thus, in this manner, the forwarding plane for the logical switch is transactionally updated in some embodiments.

(33) Instead of using the lock-based approach, other embodiments define a configuration data tuple that the controller uses to explicitly inform the MHFE that the control plane configuration data is being updated and hence might not provide a consistent view of a logical forwarding element or a logical network. For instance, instead of using the lock-based approach, some embodiments change the hardware VTEP schema to add a Boolean column to the Logical_Switch table to indicate whether the logical switch's information in the hardware VTEP database is consistent from the controller's point of view. This column is referred to below as the state_consistent column. FIG. 2 illustrates such a modified Logical_Switch table **200** with a state_consistent column **205**.

(34) In some embodiments, the controller initializes this column to False while creating a logical switch. The controller is free to update the vlan_bindings column of the Physical_Port table at any point for this logical switch. After supplying all the data tuples for the logical switch (e.g., all the data tuples to the Ucast_Macs_Remote table, Mcast_Macs_Remote table, and Physical_Port table), the controller updates the state_consistent column to True. Once that happens, the MHFE database client modifies the MHFE's forwarding tables in the data plane based on the supplied logical switch configuration tuples.

(35) When the controller wants to modify the MHFE's control plane configuration for an existing logical switch, the controller first changes the state_consistent column **205** to False. The MHFE agent monitors this column. When the column becomes False, the MHFE's forwarding records in the data plane continue to work in its last transactional state. The controller makes the changes to the MHFE's hardware VTEP database, and then changes the state_consistent column to True. At this point, the MHFE's database client applies all changes to its forwarding records in the data planes.

(36) The above-described transactional control techniques will be further described below by reference to FIGS. 5-9. However, before describing these figures, the network control system of some embodiments will be further described below by reference to FIGS. 3 and 4.

(37) FIG. 3 illustrates a network control system **300** that implements the MHFE management methods of some embodiments of the invention. In this system, the MHFE is a TOR switch **302** that communicatively couples data compute nodes (e.g., VMs, containers, etc.) that execute on host computers in a shared public network, with standalone or virtualized servers in a private network. In this example, the public network includes two racks **350** and **352** that include a plurality of VMs that execute on shared host computers. The two racks have two TORs **306** and **308**. The private network includes a rack **354** that includes virtualized and non-virtualized servers. One of ordinary skill will realize that in other examples the data compute nodes in the private network do not reside in a rack or in a single rack. Also, in some embodiments, the private and public network compute nodes can reside in the same datacenter (i.e., same physical location), or they can reside in different datacenters (i.e., in different locations).

(38) The TOR **302** implements a logical switch to connect the servers in the private network (i.e., in the rack **354**) to the VMs in the public network. As shown, the servers of the public racks **350** and **352** execute SFEs (e.g., a software switch and/or router) in addition to the VMs. These SFEs are configured by a cluster of physical controllers **330** and a cluster of logical controllers **335** to form multiple logical networks, each with one or more logical switches and routers. The logical networks isolate the data message communication of the different sets of VMs from each other, in order to allow the different sets of VMs for different logical networks to operate securely on the same and/or different hosts.

(39) In some embodiments, the host computers in the same public network rack or in two different public network racks can connect to one another through one or more tunnels **312** that allow the LFEs of the logical networks to be formed as logical overlay forwarding elements. The tunnel headers in some embodiments include logical network identifiers (e.g., VNIs) that are needed to uniquely identify the LFEs. Different types of tunnels can be used in different embodiments. Examples of such tunnels include STT (stateless transport tunnels), GRE (Generic Routing Encapsulation) tunnels, VXLAN tunnels, Geneve tunnels, etc. Tunnels can often be viewed as point-to-point logical wire connections between their endpoints (e.g., between a host and a TOR, or between two hosts) because packets inside the tunnel headers are transparent to the intervening network fabric (e.g., intervening switches, routers, etc.).

(40) In this environment, the network controllers **330** and **335** configure the TOR **302** to become part of a logical network formed by SFEs in the public network rack(s). These controllers, on the other hand, do not configure the TORs **306** and **308** to be part of the logical networks. These TORs **306** and **308** are treated as intervening network fabric. As shown, the TOR **302** connects to the host computers in the public network racks **350** and **352** through multiple logical overlay tunnels **314** for carrying the logical network identifiers for the logical network and for isolating the logical network data messages from intervening public network fabric (e.g., from TORs **306** and **308** on the racks **350** and **352**). By incorporating this TOR **302** into a logical network (e.g., into a logical switch for a logical network), the data messages from the VMs of the logical network can be directed to the ports of the TOR **302** for forwarding to DCNs (e.g., VMs and servers) in the private network rack **354**.

(41) The logical controllers generate data to define the logical forwarding elements, while the physical controllers distribute the generated data to the TOR **302** and SFEs. The number of logical controllers can be different than the number of logical networks as one logical controller can generate data for multiple logical networks. The generated data is used to configure the SFEs and TOR **302** to implement the logical forwarding elements. In some embodiments, the generated data is transformed into physical data by the physical controllers **330**, local controllers (not shown) executing on the hosts, and/or by a module operating on the TOR **302**, before this data is supplied to the forwarding plane of the SFEs and/or TOR **302**. For instance, before distributing the data generated by the logical controller, a physical controller in some embodiments converts the data into another format, e.g., into (1) physical control plane data for the TOR **302** and/or SFEs, or (2) into a format that a TOR module or host local controller can further process to produce physical control plane data.

(42) The number of physical controllers can be different than the number of managed TORs or SFEs as one physical controller typically distributes data to multiple managed TORs or SFEs. Also, in some embodiments, only one physical controller is the master controller for supplying data to a set of managed forwarding elements (e.g., SFEs or MHFEs) to configure the managed forwarding elements to facilitate the creation of LFEs. At any given time, only the master physical controller can provide data to its managed forwarding elements. In some embodiments, each forwarding element's master physical controller can have another physical controller that operates as a slave physical controller that serves as a backup (e.g., a hot standby backup) to the master physical controller in case the master controller fails.

(43) In some embodiments, one controller can operate as both a logical controller and a physical controller. Each controller in some embodiments is a separate software process, and one computing device can execute two controller processes, where one controller process is a logical controller and another controller process is a physical controller. To communicate with the managed TORs, each physical controller has a TOR agent **340** to communicate with the TORs for which the physical controller is the master controller (i.e., the primary controller for communicating with the TORs). In some embodiments, the managed TORs and TOR agents communicate with each other by using the OVSdb protocol. In some embodiments, the TOR agents employ the transactional

boundary controls of some embodiments to ensure that they provide configuration data in a transactional manner to create LFEs or to update LFEs. In some embodiments, the controllers (e.g., the logical and physical controllers **330** and **335**) communicate through RPC (remote procedure call) channels.

(44) In some embodiments, the network controller cluster defines each logical network by configuring the software and hardware forwarding elements (e.g., TOR switches). To configure such switches, some embodiments implement database servers on software and hardware forwarding elements. The network controller cluster then communicates with the database servers to provide data for configuring these software and hardware forwarding elements to implement logical forwarding elements.

(45) FIG. 4 illustrates how a network controller cluster **400** communicates with a managed software switch **405** and a managed hardware switch **410** (e.g., a TOR) in some embodiments. In this example, each of the switches has an OVSdb server **415** and **420** with which the network controller cluster **400** communicates by using the OVSdb protocol. The software switch **405** is an OVS (Open Virtual Switch) switch that executes on a host computer **425**. As shown, the software switch **405** has a database server **415**, an OVS database **430**, an OpenFlow agent **432**, and a forwarding module **434**. In the discussion below, the flow agent **432** may be referred to as an OVS daemon, and the forwarding module **434** may be referred to as a kernel module. As further shown, the hardware switch **410** includes a database server **420**, an OVS database **438**, a software stack **440**, a switch ASIC **444**, ingress and egress ports **446** and **448**, and forwarding tables **450**. The software stack **440** has a database client **452**.

(46) The network controller cluster **400** also has OVSdb clients **460** to interact with the OVSdb servers **415** and **420**. In some embodiments, the network controller cluster **400** includes logical controllers **335** and physical controllers **330** with a TOR agent **340**. In these embodiments, the physical controller **330** has an OVSdb client to interact with the OVSdb server **415** on the software switch **405**, and this controller's TOR agent has another OVSdb client to interact with the OVSdb server **420** on the hardware switch **410**.

(47) As shown, the network controller cluster **400** exchanges management data with the OVSdb server **415** of the software switch **405** by using OVSdb protocol, while exchanging configuration data with the OVS daemon **432** of the software switch by using the OpenFlow protocol. The network controller cluster **400** exchanges management data and forwarding state with the hardware switch **410** by using the OVSdb protocol.

(48) In some embodiments, the host **425** includes hardware, a hypervisor, and one or more virtual machines (VMs). The hardware may include typical computer hardware, such as processing units, volatile memory (e.g., random access memory (RAM)), nonvolatile memory (e.g., hard disc drives, optical discs, etc.), network adapters, video adapters, or any other type of computer hardware. The hardware can also include one or more NICs (network interface controllers).

(49) A hypervisor is a software abstraction layer that can run on top of the hardware of the host **425**. There are different types of hypervisors, namely Type 1 (bare metal), which runs directly on the hardware of the host, and Type 2 (hosted), which run on top of the host's operating system. The hypervisor handles various management tasks, such as memory management, processor scheduling, or any other operations for controlling the execution of the VMs. Moreover, the hypervisor communicates with the VMs to achieve various operations (e.g., setting priorities). In some embodiments, the hypervisor is a Xen hypervisor while, in other embodiments, the hypervisor may be any other type of hypervisor for providing hardware virtualization of the hardware on the host **425**.

(50) In some embodiments, the software switch **405** runs on a VM. The VM can be a unique virtual machine, which includes a modified Linux kernel (e.g., to include the OVS kernel module **434**). The VM of such embodiments is responsible for managing and controlling other VMs running on the hypervisor. In some embodiments, the VM includes a user space and the OVS daemon runs as a

background process in the user space.

(51) The OVS daemon **432** is a component of the software switch **405** that makes switching decisions. On the other hand, the kernel module **434** receives the switching decisions, caches them, and uses them subsequently to process packets. For instance, when a packet comes in, the kernel module **434** first checks a datapath cache to find a matching flow entry. If no matching entry is found, the control is shifted to the OVS daemon **432**. The OVS daemon **432** examines one or more flow tables to generate a flow to push down to the kernel module **434**. In this manner, when any subsequent packet is received, the kernel module **434** can quickly process the packet using the cached flow entry. The kernel module **434** provides a fast path to process each packet.

(52) Network controller **400** uses the OpenFlow protocol to inspect and modify a set of one or more flow tables managed by the OVS daemon **432**. The network controller cluster **400** computes flows and pushes them to the software switch **405** through this OpenFlow channel. The network controller communicates with the database server **415** of the software switch **405** by using the database protocol. Through these communications, the network controller can push configuration data for creating and managing overlay tunnels to transport nodes. The network controller might also use OVSdb protocol for discovery purposes (e.g., discover which virtual machines are hosted at the hypervisor). The OVS daemon **432** also communicates with the database server **415** to access management data (e.g., bridge information, virtual interfaces information) stored in the database **430**.

(53) Unlike its communication with the software switch **405**, the network controller **400** communicates with the hardware switch **410** by just using the OVSdb protocol. The database protocol is essentially used to control the hardware switch **410**. Through the database channel, the network controller reads the configurations from the hardware switch (e.g., an inventory of its physical ports) and sends management data to the hardware switch. For example, the network controller **400** can send instructions to the hardware switch to create tunnel ports for a logical switch. Also, when the network controller exchanges forwarding state (e.g., L2 and/or L3 forwarding state) with the hardware switch **410**, the network controller can instruct the hardware switch **410** to program its forwarding table using the database protocol.

(54) The hardware switch's ingress ports **446** are a set of ports through which the hardware switch **410** receives network data. The ingress ports **446** may include different numbers of ingress ports in different embodiments. As shown, the ingress ports **446** receives network data that is external to the switch **410**. Packets received through the ingress ports are processed by the switch's ASIC **444**.

(55) The switch ASIC is a component, which is specifically designed to support in-hardware forwarding. That is, it is primarily designed to quickly forward packets. To simplify the description, only one switching ASIC is shown. However, one of ordinary skill in the art would understand that the hardware switch **410** could include a number of ASICs that operate in conjunctions with one another to forward packets.

(56) The ASIC **444** processes the packets that it receives by using the flow entries in its forwarding tables **450**. In some embodiments, the forwarding tables **450** store active flow tables and/or flow entries that are used to determine operations for making switching decisions. In this example, each flow entry includes a qualifier and an action. The qualifier defines a set of fields to match against a set of packet header fields. As shown, the flow entries are stored in memory. The memory can be random access memory (RAM) or some other type of memory such as Content Addressable Memory (CAM) or Ternary Content Addressable Memory (TCAM). For example, a vendor may design their Layer 2 switches with CAM for performing Layer 2 switching and/or with TCAM for performing Quality of Service (QoS) functions. The switch architecture may support the ability to perform multiple lookups into multiple distinct CAM and/or TCAM regions in parallel. The CAM and TCAM are examples of switching ASICs that some vendors' switches leverage for line-speed fast switching.

(57) After processing the packet, the ASIC **444** supplies the packet to one of the egress ports **448**.

The egress ports **448** represent a set of ports through which the switching element **410** sends network data. The egress ports **448** may include different numbers of egress ports in different embodiments. In some embodiments, some or all of the egress ports **448** may overlap with some or all of the ingress ports **446**. The ports **446** and **448** along with the ASIC **444** and forwarding tables **450** compose the data plane datapath of the hardware switch **410**. The flow entries in the forwarding tables **450** represent the data plane records of the hardware switch **410**, while the database records in the OVS database **438** represent the control plane records of the hardware switch **410**.

(58) The OVSdb server **420** controls access to the database **438**. Through this server **420**, the database client **452** accesses the database **438** to read and write data. In addition, through the OVSdb server **420**, the OVSdb client **460** on the network controller **400** accesses the database **438** to read and write data (e.g., management data and forwarding state). In some embodiments, the database server **420** may send a notification to one database client (e.g., on the switch end) if the other database client (e.g., on the network controlled end) updates a table or a subset of a table of the database **438**. In some embodiments, the database protocol specifies a monitor call, which is a request, sent from a database client (**452** or **460**) to the database server **420**, to monitor one or more columns of a table and receive updates when there is an update to the one or more columns (e.g., a new row value, an update to an existing row value, etc.).

(59) For example, when the client **460** on the network controller **400** makes an update to database **438** through the OVSdb server **420**, the OVSdb server **420** in turn generates a notification for the hardware switch's client **452**. The client **452** may then read the update, and have its associated software stack **440** program the forwarding tables **450** of the switch ASIC **444**. Another example is when the database client **452** on the switch's software stack **440** updates the database **438** (through the server **420**) with MAC addresses of a machine that is connected to its port. This would in turn cause the database server **420** to send a notification regarding the update to the client **460** on the network controller **400**.

(60) In some embodiments, the database server **420** does not notify the OVSdb client **452** of the creation of a logical forwarding element data set until the `vlan_binding` column for the logical forwarding element has been defined. Alternatively, in embodiments that define a `state_consistent` tuple for the logical forwarding element, the database server **420** does not notify the OVSdb client **452** of updates to the records for a logical forwarding element after the `state_consistent` tuple is changed from True to False, until this tuple is changed back to True.

(61) In some embodiments, the database server **420** is designed to handle transactions and deal with conflicts with multiple writers (e.g., when more than one OVSdb client tries to write to the OVS database **438**). The database server **420** is also designed to provide asynchronous notifications. For example, when there is an update to a database table, the database server **420** sends a notification regarding an update to a client (e.g., executing on a network controller or on the hardware switch). In some embodiments, the database server **420** defines a lock for each logical forwarding element, and processes lock requests, lock steal requests and lock releases from the network controller's OVSdb client **460** and hardware OVSdb client **452**.

(62) The switch software stack **440** represents several programs that operate on the hardware switch **410**. The software stack **440** can include a variety of different programs to configure and manage the switch **410**. This can include management that is in and outside of the scope of the network controller cluster. For instance, the software stack **440** may include a program to update its firmware, modify switch settings (e.g., its administrative password), and/or reset the switch. The software stack **440** is vendor specific, which means that it can change from one hardware-switch vendor to another hardware-switch vendor. Accordingly, different vendors might provide different features that are represented by their corresponding software stack **440**.

(63) The software stack **440** includes at least one module to program the switch ASIC **444** and update the forwarding plane records based on control plane records retrieved from the OVS

database **438**. The software stack **440** updates the data plane records (e.g., forwarding records in the forwarding table **450**) using any number of known techniques. Different switch vendors use different techniques to update data plane records based on the retrieved control plane records.

(64) FIG. 5 illustrates a process **500** that a physical controller performs in some embodiments to supply the definition of a logical forwarding element to an MHFE in a transactional manner. In this example, the TOR agent of the physical controller performs this process to deploy a logical switch on a TOR in a transactional manner. Many of the embodiments described below by reference to FIGS. **6-9** are also described by reference to operations of TORs and the creation of logical switches. However, one of ordinary skill will realize that other embodiments perform analogous processes to define other types of logical forwarding elements and/or to deploy logical forwarding elements on different types of MHFEs.

(65) As shown, the process **500** initially receives (at **505**) the definition of a logical switch that needs to be deployed on a TOR. Next, the process **500** selects (at **510**) a table to update in the TOR's OVS database **438** other than the Physical_Port table. Examples of such tables include the Ucast_Macs_Remote table, the Mcast_Macs_Remote table and the Logical_Switch table. At **515**, the process **500** sends to the TOR's OVSdb server **420** one or more packets that contain data for defining one or more records in the selected table (i.e., the table selected at **510**) for the logical switch. Examples of such records include the logical switch record in the Logical_Switch table **115**, MAC records in the Ucast table **105**, and service node records in the Mcast table **110**, as described above by reference to FIG. 1.

(66) Next, at **520**, the process **500** determines whether it has pushed the data for all the OVS database tables except the Physical_Port table to the TOR's OVSdb server **420**. If not, the process **500** returns to **510** to select another OVSdb table other than the Physical_Port table, and then to **515** to push data records for this newly selected table to the TOR's OVSdb server **420**. For some embodiments, the illustration of the determination operation **520** is conceptual as in these embodiments the TOR agent selects the OVSdb tables according to a fixed order, in which the Physical_Port table is last.

(67) Once the process **500** determines (at **520**) that it has processed all the OVSdb tables except for the Physical_Port table, the process determines (at **525**) whether it has received confirmation from the OVSdb server **420** that it has created and/or populated the appropriate records in the OVS database **438** with the data that the process **500** send at **515**. If not, the process **500** remains at **525** until it has received the required confirmations. In some cases, the process **500** resends data for an OVSdb table when it does not receive the required confirmation for this table.

(68) Once the process **500** determines (at **525**) that it has received confirmations for all the OVSdb tables for which it has pushed data to the TOR, the process transitions to **530**. In other embodiments, the process **500** does not check (at **525**) to see whether it has received confirmations from the OVSdb server **420** that it has created and/or populated the appropriate records in the OVS database **438** based on the data provided by the process **500**. Instead, in these embodiments, the process **500** transitions to **530** when it determines (at **520**) that it has the data for the Physical_Port table to process.

(69) At **530**, the process **500** sends the VLAN binding tuple or tuples for the Physical_Port table to the OVSdb server **420**. The process **500** remains at **530** until it receives confirmation from the OVSdb server **420** that it has processed the VLAN binding tuples. Once the process **500** receives this confirmation, it ends. Again, the process **500** in some embodiments does not wait for confirmation from the OVSdb server **420**, and just ends after it sends the VLAN binding tuple or tuples for the Physical_Port table.

(70) The process **500** pushes (at **530**) the VLAN binding tuples last in order to prevent the hardware switch's OVSdb client **452** from updating the forwarding tables **450** until all of the logical switch's records have been specified in the OVS database **438**. In some embodiments, the OVSdb server **420** does not notify the OVSdb client **452** of changes to the OVS database **438** relating to a

logical switch until the VLAN binding column is populated. Hence, by delaying pushing the VLAN binding tuples, the process **500** can prevent the OVSdb client **452** from updating the forwarding plane until the process has pushed all the control plane records for the logical switch to the hardware switch.

(71) FIG. **6** illustrates a process **600** that a hardware switch's OVSdb server **420** performs to define a logical switch on the hardware switch. As shown, this process starts when the OVSdb server **420** receives (at **605**) data tuples from the process **500** of the network controller (e.g., TOR agent) for the logical switch. These data tuples include tuples for the OVSdb tables that need to have records for the logical switch to be defined on the hardware switch. These tables include the

Ucast_Macs_Remote table **105**, the Mcast_Macs_Remote table **105** and the Logical_Switch table **115**. The received data tuples (at **605**) do not include data tuples for the Physical_Port table **120**.

(72) At **605**, the process **600** creates the records and/or populates previously created records in these tables based on the received data tuples. At **610**, the process **600** then sends confirmation to the process **500** of the network controller of its processing of the received data tuples (i.e., of its creation or population of the data records). To avoid obscuring the description of the process **600** with unnecessary detail, FIG. **6** illustrates that the process **600** supplies confirmations for all non-Physical_Port table tuples at once after receiving all of these data tuples. One of ordinary skill will realize that in some embodiments the process **600** supplies a confirmation for each set of such tuples (e.g., for each set of tuples for a table) after processing that set of tuples.

(73) Once the process **600** has provided (at **610**) confirmation of processing of all the received data tuples, the process receives (at **615**) the VLAN-binding tuple(s) for the logical switch. The process **600** then records the received VLAN-binding tuple(s) in the vlan_bindings column of the Physical_Port table. As populated, each VLAN binding tuple in the vlan_bindings column of the Physical_Port table binds a specific VLAN on a specific physical port of the hardware switch to the logical switch. In some embodiments, the process **600** receives only one VLAN-binding tuple for each logical switch for the hardware switch, because in these embodiments each logical switch can bind to at most one VLAN binding of a physical port. In other embodiments, however, the process can receive more than one VLAN-binding tuple for each logical switch for the hardware switch, because in these embodiments each logical switch can bind to more than one VLAN binding of a physical port or to more than one port.

(74) After storing (at **615**) the received VLAN-binding tuple(s) in the Physical_Port table, the process **600** creates (at **620**) a lock for the logical switch. In the embodiments that do not utilize a lock for modifying a logical switch's configuration on a hardware switch, no lock is defined at **620**. On the other hand, in the embodiments in which the controller needs a lock to modify the logical switch's records, the controller can steal the logical switch's lock, in order to prevent the hardware switch's OVSdb client from modifying the data plane records based on control plane records that the controller has not been able to fully update. In some embodiments, the process **600** creates a lock for the logical switch at the request of the process **500** of the network controller. Specifically, in these embodiments, the configuration data provided by the network controller includes a request to create a lock for the logical switch. In other embodiments, the process **600** of the OVSdb server is pre-configured to create locks for logical switches, and hence does not need to be directed to do this by the network controller.

(75) Next, at **625**, the process **600** sends a confirmation to the TOR agent that it has processed the VLAN-binding tuple(s), and created the logical switch. In some embodiments, the process **600** also provides (at **625**) a logical switch name and/or this switch's lock name or attribute to the process **500** so that the network controller (e.g., the TOR agent) will have a handle for requesting the lock for a subsequent logical switch update. In other embodiments, the process **600** does not provide to the process **500** the logical switch's name/attribute as the controller operates on an assumption that the logical switch has a lock.

(76) After providing (at **625**) the confirmation to the process **500**, the process **600** notifies (at **630**)

the TOR's OVSdb client **452** of the creation of records in the OVS database **438** for the new logical switch. At this point, the OVSdb client **452** begins to access the OVS database **438** through the OVSdb server **420**, in order to update the TOR's forwarding plane records (e.g., forwarding tables **450**). After **630**, the process ends.

(77) As described above, the MHFEs in some embodiments contain configuration databases and execute database servers with which the network controllers can interact through a database protocol. The database protocol of some embodiments provides a lock operation to lock or unlock an LFE's records in an MHFE's configuration database. The lock operation allows the network controller to steal the lock in order to prevent an MHFE database client from obtaining the lock and reading the LFE's records in the configuration database, while the controller writes to the configuration database. The lock feature resolves conflicts by making the MHFE agent receive the lock, and hence receive permission, before it can read the LFE's records in the database.

(78) More specifically, in some embodiments, a controller must obtain a lock for an LFE from the MHFE's database server before the controller can modify the LFE's records in the database. In some embodiments, the controller obtains the lock through a steal request in order (1) to be able to get the lock irrespective of whether the MHFE database client currently has the lock and (2) to be able to block the MHFE database client from getting back the lock. Similarly, an MHFE database client must obtain a lock for an LFE from the MHFE's database server before this client can read the LFE's records in the database. When the controller has stolen the lock, the MHFE database client cannot obtain the lock through a lock request until the controller releases the lock. After receiving a lock release from the controller, the database server provides the lock to the MHFE agent so that it can then read the LFE's records in the database.

(79) FIG. 7 illustrate a process **700** that the TOR agent, OVSdb server **420** and the OVSdb client **452** perform in some embodiments during and after a period in which the TOR agent modifies the records of a logical switch in the OVS database **438**. As shown, the process **700** starts when the TOR agent receives (at **702**) updates to one or more of the logical switch's data tuples from the logical controller **335** for the logical switch.

(80) Next, at **704**, the OVSdb client **460** of the TOR agent **340** provides a lock steal request for the logical switch to the OVSdb server **420** of the TOR. In some embodiments, the OVSdb client **460** first provides a lock request to the OVSdb server **420**, and if it is told that the lock currently resides with the TOR's OVSdb client **452**, then provides a lock steal request to the OVSdb server **420**. In other embodiments, however, the OVSdb client **460** just provides a lock steal request without first providing a lock request and having this lock request rejected, as illustrated in FIG. 7.

(81) After receiving (at **706**) this request, the OVSdb server **420** sends (at **708**) a notification to the TOR's OVSdb client **452** that it has lost the lock. As shown, prior to receiving this notification (at **736**), the TOR's OVSdb client **452** in this example had obtained the lock (at **734**) so that it could access the OVS database **438**.

(82) After notifying the OVSdb client **452** that it no longer has the lock for the logical switch, the OVSdb server **420** provides (at **708**) this lock to the TOR agent's OVSdb client **460**. Upon receiving (at **710**) the lock for the logical switch, the OVSdb client **460** sends (at **712**) one or more logical-switch update data tuples to the OVSdb server **420**. Based on the update data tuples, the OVSdb server **420** modifies (at **714**) one or more logical switch's records in the OVS database **438**.

(83) Once the OVSdb server **420** has received and processed all update data tuples from the TOR agent's OVSdb client **460**, the OVSdb server **420** provides (at **714**) a confirmation to the OVSdb client **460** that it has processed all of the update data tuples. After receiving (at **712**) this confirmation, the TOR agent's OVSdb client **460** sends (at **716**) a lock release command to the OVSdb server **420** so that it can relinquish the logical switch's lock. At **718**, the OVSdb server **420** receives the lock release.

(84) During the period in which the TOR agent's OVSdb client **460** has acquired the lock for the logical switch, the TOR's OVSdb server **420** does not respond to any lock request from the TOR's

OV Sdb client **452** for the logical switch. FIG. 7 pictorially illustrates a lock request **738** by the OV Sdb client **452**. As shown, the OV Sdb server **420** does not respond to this lock request while the TOR agent has acquired the lock for the logical switch. In some embodiments, the OV Sdb server **420** affirmatively rejects a lock request from the OV Sdb client **452** while the TOR agent has acquired the logical switch's lock, instead of just not responding to such a request.

(85) After the OV Sdb server receives (at **718**) the lock release from the TOR agent, the OV Sdb server supplies (at **720**) a notification to the TOR's OV Sdb client **452** of the changes to the logical switch's records in the OVS database **438**. At **720**, the OV Sdb server also provides the lock for the logical switch to the TOR's OV Sdb client **752**. The OV Sdb client **752** receives (at **722**) the lock release and change notification, and in response, directs (at **724**) the OV Sdb server to supply the updated logical switch records to it.

(86) While the example illustrated in FIG. 7 shows the OV Sdb server **420** providing both the lock and change notification at **720**, the OV Sdb server **420** of some embodiments only provides the lock at **720**. In these embodiments, when the OV Sdb client **452** automatically checks for changes to the logical switch records when it obtains the lock. Also, in some embodiments, the OV Sdb server **420** just provides the change notification at **720**. In response to this notification, the OV Sdb client **452** then asks the OV Sdb server for the lock before asking the OV Sdb server for the logical switch's records.

(87) In the example illustrated in FIG. 7, the OV Sdb server receives (at **726**) the request for the logical switch's records, and provides (at **728**) the updated logical switch records to the TOR's OV Sdb client **752**. At **730**, the OV Sdb client then modifies the forwarding plane records (e.g., the records in the forwarding tables **450**) based on the received logical switch records. In some embodiments, the TOR's OV Sdb client **752** provides a confirmation of recording these updates to the OV Sdb server. In other embodiments, the client does not provide such a confirmation.

(88) After updating the forwarding plane records, the OV Sdb client relinquishes the lock and the process **700** then ends. In some embodiments, the OV Sdb client does not relinquish the lock for the logical switch even after updating the forwarding plane records for the logical switch based on a control plane update for the logical switch. In these embodiments, the OV Sdb server affirmatively takes away the lock from the TOR's OV Sdb client **452** once the TOR agent's OV Sdb client **460** wants to modify the logical switch's control plane records.

(89) FIG. 7 illustrates the OV Sdb server **420** as sending (at **720**) one update notification to the TOR's OV Sdb client **452** once the TOR agent has completed its control plane update for the logical switch. In other embodiments, the OV Sdb server **420** or some other notification process associated with the OVS database **438** notifies the TOR's OV Sdb client **452** with multiple update notifications while the TOR agent is updating multiple data tuple sets associated with the logical switch. In these embodiments, the TOR's OV Sdb client **452** repeatedly attempts to obtain the logical switch's lock in response to the various different notifications, but these attempts are unsuccessful while the TOR agent's OV Sdb client **460** has not released the lock.

(90) In some embodiments, the OV Sdb server **420** provides a logical switch's lock to the TOR agent's OV Sdb client **460** anytime that this client asks for the lock. Other embodiments, however, place some restriction on the TOR agent's OV Sdb client **460** obtaining the switch's lock. For instance, after the OV Sdb client **460** has supplied one set of updates to the logical switch's control plane records and has relinquished the logical switch's lock, the OV Sdb server **420** does not allow this client **460** to reacquire the logical switch's lock while the TOR OV Sdb client **452** is updating the TOR's forwarding plane records in view of the supplied control plane updates. In these embodiments, the OV Sdb server **420** only allows the TOR agent's OV Sdb client **460** to reacquire the logical switch's lock after receiving a confirmation from the TOR's OV Sdb client **452** that it has processed all of the updates supplied by the TOR agent.

(91) FIG. 8 illustrates a process **800** that a TOR agent performs in some embodiments to define a logical switch or to modify a logical switch definition in a transactional manner for a TOR by using

a state_consistent tuple **205** for the logical switch. As shown, this process initially (at **805**) receives the definition of a logical switch that needs to be deployed on a TOR, or receives a modification to the definition of a previously defined logical switch that has already been deployed on the TOR. When the logical switch is a new logical switch that is being deployed on the TOR, the process sends (at **810**) one or more packets to the TOR's OVSdb server **420** to create a record for this logical switch in the logical switch table **200**, and to set the state_consistent value in this record for the logical switch to False.

(92) On the other hand, when the logical switch was previously deployed on the TOR, the process sends (at **810**) one or more packets to the TOR's OVSdb server **420** to direct it to change the state_consistent value of the logical switch to False in the logical switch table **200**. Changing this state to False prevents the TOR's OVSdb client **452** from accessing the OVS database records for this logical switch. In some embodiments, when this state is False, the TOR's OVSdb server **420** ignores or rejects the TOR's OVSdb client's **452** requests for access to the records in the OVS database **438** that relate to the logical switch.

(93) At **815**, the process **800** selects a table to update in the TOR's OVS database other than the Logical_Switch table. Examples of such tables include the Ucast_Macs_Remote table, the Mcast_Macs_Remote table and the Physical_Port table. At **820**, the process sends to the TOR's OVSdb server **420** one or more packets that contain data for defining one or more records in the selected table (i.e., the table selected at **815**) for the logical switch. Examples of such records include the logical switch record in the VLAN bindings for the Physical_Port table **120**, MAC records in the Ucast table **105**, and service node records in the Mcast table **110**, as described above by reference to FIG. 1.

(94) Next, at **825**, the process **800** determines whether it has pushed to the TOR's OVSdb server **420** the data for all the OVS database tables other than the Logical_Switch table. If not, the process **800** returns to **815** to select another OVSdb table other than the Logical_Switch table, and then to **820** to push data records for this newly selected table to the TOR's OVSdb server **420**. For some embodiments, the illustration of the determination operation **820** is conceptual as in these embodiments the TOR agent selects the OVSdb tables according to a fixed order, in which the Logical_Switch table is last.

(95) Once the process **800** determines (at **825**) that it has processed all the OVSdb tables other than the Logical_Switch table, the process **800** determines (at **830**) whether it has received confirmation from the TOR's OVSdb server **420** that it has created and/or populated the appropriate records in the OVS database **438**. If so, the process **800** transitions to **835**. If not, the process **800** remains at **830** until it has received the required confirmations. In some cases, the process **800** resends data for an OVSdb table when it does not receive the required confirmation for this table. Also, in some cases, the process **800** does not wait for confirmation from the OVSdb server **420** that it has processed the previously supplied data tuples. In these embodiments, the process **800** simply transitions from **825** to **835** when it determines that it has processed all the OVSdb tables other than the Logical_Switch table.

(96) At **835**, the process directs (at **835**) the TOR's OVSdb server **420** to change the state_consistent tuple in the Logical_Switch table **200** to True. When the logical switch's record in this table has to be modified, the process **800** in some embodiments also sends (at **835**) other data tuples to change other parameters in the Logical_Switch table for the logical switch. In some embodiments, the process **800** sends such other data tuple(s) before changing the state_consistent tuple to True. In other embodiments, the process **800** sends the other logical_switch tuple(s) for changing a pre-existing logical switch's record in the table **200** at **810**, instead of **835**.

(97) The process **800** remains at **835** until it receives confirmation from the TOR that it has changed the state_consistent tuple to True for the logical switch. Once the process receives this confirmation, it ends. The process **800** changes the state_consistent tuple to True last in order to prevent the OVSdb client **452** from updating the forwarding tables **450** until all of the logical

switch's records have been specified/modified in the OVS database **438**. In some embodiments, the OVSdb server **420** does not notify the OVSdb client **452** of changes to the OVS database **438** relating to a logical switch until the state_consistent tuple is changed from False to True. In other embodiments, the OVSdb client **452** receives notifications of changes to the OVS database **438** relating to a logical switch even when the state_consistent tuple is False. However, while this value is False, the OVSdb server **420** ignores or rejects all requests by the OVSdb client **452** for access to records in the OVS database **438** relating to the logical switch. In either of these approaches, by delaying pushing the state_consistent tuple, the process **800** can prevent the OVSdb client **452** from updating the forwarding plane until the process has pushed all the control plane records or modifications for the logical switch to the TOR.

(98) FIG. **9** illustrates a process **900** that a hardware switch's OVSdb server **420** performs to create a logical switch on the hardware switch, or to update the configuration of a previously created logical switch on the hardware switch. This process starts when it OVSdb server **420** receives (at **905**) from the network controller data tuples for the logical switch to the OVSdb server **420**. When the logical switch was previously created, the process **900** initially receives (at **905**) a request from the process **800** to change the logical switch's state_consistent tuple to False. In response, the process **900** (at **910**) changes this tuple to False in the Logical_Switch table **200**, sends a notification to the hardware switch's OVSdb client **452** of this changed state, and then sends a confirmation of this change to the process **800** (i.e., to the TOR agent of the network controller).

(99) On the other hand, when the logical switch was not previously created, the process **900** receives (at **905**) one or more data tuples to create a record for this logical switch in the logical switch table **200**, and to set the state_consistent value in this record for the logical switch to False. In response, the process **900** (at **910**) creates a record in the Logical_Switch Table **200** for the logical switch, sets the state_consistent tuple in this record to False, and then sends a confirmation of the creation of the record to the process **800** (i.e., to the TOR agent of the network controller).

(100) At **915**, the process **900** receive data tuples from the process **800** of the network controller (e.g., TOR agent) for the logical switch. These data tuples include tuples for the OVSdb tables that need to have records for the logical switch to be defined on the hardware switch. These tables include the Ucast_Macs_Remote table **105**, the Mcast_Macs_Remote table **110** and the Physical_Port table **120**. These data tuples (received at **915**) do not include data tuples for the Logical_Switch table **115**.

(101) At **915**, the process creates the records and/or populates previously created records in these tables based on the received data tuples. At **920**, the process then sends confirmation to the process **800** of the network controller of its processing of the received data tuples (i.e., of its creation or population of the data records in the non Logical_Switch tables of the OVS database **438** for the logical switch). To avoid obscuring the description of the process **900** with unnecessary detail, FIG. **9** illustrates that the process **900** supplies a confirmation for all non-Logical_Switch table tuples at once after receiving all of these data tuples. One of ordinary skill will realize that in some embodiments the process **900** supplies a confirmation for each set of such tuples (e.g., for each set of tuples for a table) after processing that set of tuples.

(102) Once the process **900** has provided confirmation of processing of all the received data tuples, the process (at **925**) receives a data tuple to change the state_consistent tuple to True and change this state to True. When the logical switch's record in this table has to be modified, the process **900** in some embodiments also receives (at **925**) other data tuples to change other parameters in the Logical_Switch table for the logical switch. In some embodiments, the process **900** receives such other data tuple(s) and processes these data tuples (e.g., changes records in the database tables based on these tuples) before changing the state_consistent tuple back to True. In other embodiments, the process **900** receives the other logical switch tuple(s) for changing a pre-existing logical switch's record in the table **200** at **905** instead of **925**, and it changes the Logical_Switch table **200** tuple(s) at **910** instead of **925**.

(103) After storing (at **915**) the received data tuple(s) in the Logical_Switch table, the process **900** sends (at **920**) a confirmation to the process **800** that it has processed the tuple(s) Logical_Switch table. When the logical switch was not previously created, this confirmation also informs the process **800** of the creation of the logical switch. In some embodiments, the process **900** also provides (at **920**) a logical switch name to the process **800** so that the network controller (e.g., the TOR agent) will have a handle for the logical switch. In other embodiments, the process **900** does not provide the logical switch's name to the process **800** as the controller has this switch's identifier.

(104) Once the process **900** has made its change(s) to the Logical_Switch table **200** at **925**, the process sends a confirmation of its processing the data tuple(s) received at **925**. After providing (at **925**) the confirmation to the process **800**, the process **900** notifies (at **930**) the TOR's OVSdb client **452** of the creation of records in the OVSdb database **438** for the new logical switch or of the change of state_consistent tuple to True for a previously created logical switch. At this point, the OVSdb client **452** can begin to access the OVS database **438** through the OVSdb server **420**, in order to update the TOR's forwarding plane records (e.g., forwarding tables **450**). After **925**, the process ends.

(105) One of ordinary skill will realize that the approach illustrated in FIGS. **8** and **9** is just one way of using the state_consistent tuple to create/update a logical switch's records on a TOR. Other embodiments use this data tuple differently. For instances, some embodiments do not first create the logical switch's record in the Logical_Switch table and define the state_consistent tuple in this record to False, when the logical switch is being initially defined on a TOR. These embodiments just create the logical switch's record in the Logical_Switch table last. In some of these embodiments, the hardware VTEP schema is modified to allow the other tables (e.g., Ucast table **105** and Mcast table **110**) to be created before the logical switch's record is created in the Logical_Switch table **200**.

(106) Many of the above-described features and applications are implemented as software processes that are specified as a set of instructions recorded on a computer readable storage medium (also referred to as computer readable medium). When these instructions are executed by one or more processing unit(s) (e.g., one or more processors, cores of processors, or other processing units), they cause the processing unit(s) to perform the actions indicated in the instructions. Examples of computer readable media include, but are not limited to, CD-ROMs, flash drives, RAM chips, hard drives, EPROMs, etc. The computer readable media does not include carrier waves and electronic signals passing wirelessly or over wired connections.

(107) In this specification, the term “software” is meant to include firmware residing in read-only memory or applications stored in magnetic storage, which can be read into memory for processing by a processor. Also, in some embodiments, multiple software inventions can be implemented as sub-parts of a larger program while remaining distinct software inventions. In some embodiments, multiple software inventions can also be implemented as separate programs. Finally, any combination of separate programs that together implement a software invention described here is within the scope of the invention. In some embodiments, the software programs, when installed to operate on one or more electronic systems, define one or more specific machine implementations that execute and perform the operations of the software programs.

(108) FIG. **10** conceptually illustrates a computer system **1000** with which some embodiments of the invention are implemented. The computer system **1000** can be used to implement any of the above-described hosts, controllers, and managers. As such, it can be used to execute any of the above described processes. This computer system includes various types of non-transitory machine readable media and interfaces for various other types of machine readable media. Computer system **1000** includes a bus **1005**, processing unit(s) **1010**, a system memory **1025**, a read-only memory **1030**, a permanent storage device **1035**, input devices **1040**, and output devices **1045**.

(109) The bus **1005** collectively represents all system, peripheral, and chipset buses that communicatively connect the numerous internal devices of the computer system **1000**. For

instance, the bus **1005** communicatively connects the processing unit(s) **1010** with the read-only memory **1030**, the system memory **1025**, and the permanent storage device **1035**.

(110) From these various memory units, the processing unit(s) **1010** retrieve instructions to execute and data to process in order to execute the processes of the invention. The processing unit(s) may be a single processor or a multi-core processor in different embodiments. The read-only-memory (ROM) **1030** stores static data and instructions that are needed by the processing unit(s) **1010** and other modules of the computer system. The permanent storage device **1035**, on the other hand, is a read-and-write memory device. This device is a non-volatile memory unit that stores instructions and data even when the computer system **1000** is off. Some embodiments of the invention use a mass-storage device (such as a magnetic or optical disk and its corresponding disk drive) as the permanent storage device **1035**.

(111) Other embodiments use a removable storage device (such as a floppy disk, flash drive, etc.) as the permanent storage device. Like the permanent storage device **1035**, the system memory **1025** is a read-and-write memory device. However, unlike storage device **1035**, the system memory is a volatile read-and-write memory, such a random access memory. The system memory stores some of the instructions and data that the processor needs at runtime. In some embodiments, the invention's processes are stored in the system memory **1025**, the permanent storage device **1035**, and/or the read-only memory **1030**. From these various memory units, the processing unit(s) **1010** retrieve instructions to execute and data to process in order to execute the processes of some embodiments.

(112) The bus **1005** also connects to the input and output devices **1040** and **1045**. The input devices enable the user to communicate information and select commands to the computer system. The input devices **1040** include alphanumeric keyboards and pointing devices (also called "cursor control devices"). The output devices **1045** display images generated by the computer system. The output devices include printers and display devices, such as cathode ray tubes (CRT) or liquid crystal displays (LCD). Some embodiments include devices such as a touchscreen that function as both input and output devices.

(113) Finally, as shown in FIG. **10**, bus **1005** also couples computer system **1000** to a network **1065** through a network adapter (not shown). In this manner, the computer can be a part of a network of computers (such as a local area network ("LAN"), a wide area network ("WAN"), or an Intranet, or a network of networks, such as the Internet. Any or all components of computer system **1000** may be used in conjunction with the invention.

(114) Some embodiments include electronic components, such as microprocessors, storage and memory that store computer program instructions in a machine-readable or computer-readable medium (alternatively referred to as computer-readable storage media, machine-readable media, or machine-readable storage media). Some examples of such computer-readable media include RAM, ROM, read-only compact discs (CD-ROM), recordable compact discs (CD-R), rewritable compact discs (CD-RW), read-only digital versatile discs (e.g., DVD-ROM, dual-layer DVD-ROM), a variety of recordable/rewritable DVDs (e.g., DVD-RAM, DVD-RW, DVD+RW, etc.), flash memory (e.g., SD cards, mini-SD cards, micro-SD cards, etc.), magnetic and/or solid state hard drives, read-only and recordable Blu-Ray® discs, ultra density optical discs, any other optical or magnetic media, and floppy disks. The computer-readable media may store a computer program that is executable by at least one processing unit and includes sets of instructions for performing various operations. Examples of computer programs or computer code include machine code, such as is produced by a compiler, and files including higher-level code that are executed by a computer, an electronic component, or a microprocessor using an interpreter.

(115) While the above discussion primarily refers to microprocessor or multi-core processors that execute software, some embodiments are performed by one or more integrated circuits, such as application specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs). In some embodiments, such integrated circuits execute instructions that are stored on the circuit itself.

(116) As used in this specification, the terms “computer”, “server”, “processor”, and “memory” all refer to electronic or other technological devices. These terms exclude people or groups of people. For the purposes of the specification, the terms display or displaying means displaying on an electronic device. As used in this specification, the terms “computer readable medium,” “computer readable media,” and “machine readable medium” are entirely restricted to tangible, physical objects that store information in a form that is readable by a computer. These terms exclude any wireless signals, wired download signals, and any other ephemeral or transitory signals.

(117) While the invention has been described with reference to numerous specific details, one of ordinary skill in the art will recognize that the invention can be embodied in other specific forms without departing from the spirit of the invention. For instance, this specification refers throughout to computational and network environments that include virtual machines (VMs). However, virtual machines are merely one example of data compute nodes (DCNs) or data compute end nodes, also referred to as addressable nodes. DCNs may include non-virtualized physical hosts, virtual machines, containers that run on top of a host operating system without the need for a hypervisor or separate operating system, and hypervisor kernel network interface modules.

(118) VMs, in some embodiments, operate with their own guest operating systems on a host using resources of the host virtualized by virtualization software (e.g., a hypervisor, virtual machine monitor, etc.). The tenant (i.e., the owner of the VM) can choose which applications to operate on top of the guest operating system. Some containers, on the other hand, are constructs that run on top of a host operating system without the need for a hypervisor or separate guest operating system. In some embodiments, the host operating system uses name spaces to isolate the containers from each other and therefore provides operating-system level segregation of the different groups of applications that operate within different containers. This segregation is akin to the VM segregation that is offered in hypervisor-virtualized environments that virtualize system hardware, and thus can be viewed as a form of virtualization that isolates different groups of applications that operate in different containers. Such containers are more lightweight than VMs.

(119) Hypervisor kernel network interface module, in some embodiments, is a non-VM DCN that includes a network stack with a hypervisor kernel network interface and receive/transmit threads. One example of a hypervisor kernel network interface module is the vmknix module that is part of the ESXi™ hypervisor of VMware, Inc. One of ordinary skill in the art will recognize that while the specification refers to VMs, the examples given could be any type of DCNs, including physical hosts, VMs, non-VM containers, and hypervisor kernel network interface modules. In fact, the example networks could include combinations of different types of DCNs in some embodiments.

(120) A number of the figures (e.g., FIGS. 5-9) conceptually illustrate processes. The specific operations of these processes may not be performed in the exact order shown and described. The specific operations may not be performed in one continuous series of operations, and different specific operations may be performed in different embodiments. Furthermore, the process could be implemented using several sub-processes, or as part of a larger macro process. In view of the foregoing, one of ordinary skill in the art would understand that the invention is not to be limited by the foregoing illustrative details, but rather is to be defined by the appended claims.

Claims

1. A non-transitory machine readable medium storing a program for configuring a managed hardware forwarding element (MHFE) to implement a logical forwarding element (LFE) along with a plurality of other managed forwarding elements operating outside of the MHFE, and to communicatively couple the LFE with a private network, the program comprising sets of instructions for: providing a plurality of data tuples for defining the LFE on the MHFE, without providing a data tuple for binding the LFE to a port of the MHFE to ensure that all data tuples of the plurality of data tuples are received at the MHFE before the MHFE creates forwarding records

- for the LFE; after providing the plurality of data tuples for defining the LFE, receiving confirmation from the MHFE that the MHFE has processed the plurality of data tuples; and after receiving the confirmation, providing the data tuple that binds the LFE to a port of the MHFE in order to couple the LFE with the private network through the MHFE.
2. The non-transitory machine readable medium of claim 1, wherein the program further comprises a set of instructions for ordering the data tuples for the MHFE such that the data tuple for binding the LFE to the MHFE port is last.
 3. The non-transitory machine readable medium of claim 1, wherein: the MHFE comprises (i) a module that retrieves the provided data tuples and uses the provided data tuples to configure dataplane forwarding records of the MHFE to implement the LFE and (ii) a forwarding engine that uses the dataplane forwarding records to process data messages that the MHFE receives; and the module does not configure the dataplane forwarding records based on the provided data tuples until the data tuple for binding the LFE to the MHFE port is provided.
 4. The non-transitory machine readable medium of claim 1, wherein: the MHFE comprises a database with a plurality of tables for storing the provided data tuples; the database uses a hardware VTEP (VXLAN Tunnel End Point) schema and the data tuples are provided to the database on the MHFE by using an OVSdb (open vswitch database) protocol.
 5. The non-transitory machine readable medium of claim 1, wherein the MHFE port is associated with multiple VLANs (virtual local area networks), and the binding data tuple binds the LFE to a VLAN associated with the MHFE port.
 6. The non-transitory machine readable medium of claim 1, wherein the other forwarding elements that implement the LFE with the MHFE comprise a plurality of software forwarding elements that execute on host computers on which data compute nodes also execute.
 7. A method for configuring a managed hardware forwarding element (MHFE) to implement a logical forwarding element (LFE) along with a plurality of other managed forwarding elements operating outside of the MHFE, and to communicatively couple the LFE with a private network, the method comprising: providing a plurality of data tuples to the MHFE to implement the LFE along with the plurality of other managed forwarding elements operating outside of the MHFE, said plurality of data tuples provided to the MHFE without providing a data tuple for binding the LFE to a port of the MHFE; after providing the plurality of data tuples for implementing the LFE, receiving confirmation from the MHFE that the MHFE has processed the plurality of data tuples to implement the LFE on the MHFE; and after receiving the confirmation, providing the data tuple that binds the LFE to a port of the MHFE in order to couple the LFE with the private network through the MHFE.
 8. The method of claim 7 further comprising ordering the data tuples for the MHFE such that the data tuple for binding the LFE to the MHFE port is last.
 9. The method of claim 7: wherein the MHFE comprises (i) a module that retrieves the provided data tuples and uses the provided data tuples to configure dataplane forwarding records of the MHFE to implement the LFE, and (ii) a forwarding engine that uses the dataplane forwarding records to process data messages that the MHFE receives; and the module does not configure the dataplane forwarding records based on the provided data tuples until the data tuple for binding the LFE to the MHFE port is provided.
 10. The method of claim 7, wherein: the MHFE comprises a database with a plurality of tables for storing the provided data tuples; the database uses a hardware VTEP (VXLAN Tunnel End Point) schema; and the data tuples are provided to the database on the MHFE by using an OVSdb (open vswitch database) protocol.
 11. The method of claim 7, wherein the MHFE port is associated with multiple VLANs (virtual local area networks), and the binding data tuple binds the LFE to a VLAN associated with the MHFE port.
 12. The method of claim 7, wherein the other forwarding elements that implement the LFE with the

MHFE comprise a plurality of software forwarding elements that execute on host computers on which data compute nodes also execute.

13. A non-transitory machine readable medium storing a program for configuring a managed hardware forwarding element (MHFE) to implement a logical forwarding element (LFE) along with a plurality of other managed forwarding elements operating outside of the MHFE, and to communicatively couple the LFE with a private network, the program comprising sets of instructions for: providing a first set of data tuples to define the LFE on the MHFE and to define a state of the LFE as incomplete; providing a second set of data tuples to complete the definition of the LFE on the MHFE; after providing the first and second sets of data tuples, receiving confirmation that the first and second sets of data tuples have been processed by the MHFE; and after receiving confirmation, providing a third set of data tuples to change the LFE state to complete.

14. The non-transitory machine readable medium of claim 13, wherein the MHFE comprises (i) a module that retrieves the provided data tuples and uses the provided data tuples to configure dataplane forwarding records of the MHFE to implement the LFE and (ii) a forwarding engine that uses the dataplane forwarding records to process data messages that the MHFE receives; and the module does not configure the dataplane forwarding records based on the provided data tuples until the state of the LFE is set to complete.

15. The non-transitory machine readable medium of claim 13, wherein the MHFE comprises a database with a plurality of tables for storing the provided data tuples.

16. The non-transitory machine readable medium of claim 15, wherein the database uses a hardware VTEP (VXLAN Tunnel End Point) schema, and the data tuples are provided to the database on the MHFE by using an OVSdb (open vswitch database) protocol.

17. The non-transitory machine readable medium of claim 16, wherein: the LFE is a logical switch; the hardware VTEP schema specifies a logical switch table that comprises a record for the logical switch; and the logical switch record includes a Boolean field for specifying whether the state of the logical switch is complete.
