

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388757
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Danivas; Vishwas et al.

Systems and methods for using programmable policer circuits for network flow policing

Abstract

The packet processing chip of a networking device includes a packet processing pipeline circuit and a programmable policer circuit. A single programmable policer circuit may use policing policy identifiers and an aggregated token bucket to police multiple network flows. A policing policy identifier may govern which policing policy is used for a network packet. Each network flow may have a flow table entry that includes a policing policy identifier and a programmable policer circuit identifier. After reading the flow table entry for processing a network packet, the packet processing pipeline circuit may send data including the policing policy identifier to the programmable policer circuit which returns a policing decision in accordance with the policing policy and the state of the aggregated token bucket. Different network flows may use the same aggregated token bucket and different policing policies to thereby implement strict priority of some network flows over others.

Inventors: Danivas; Vishwas (Santa Clara, CA), Chu; Kit Chiu (Fremont, CA), Kotha; Murty (San Jose, CA)

Applicant: Advanced Micro Devices, Inc. (Santa Clara, CA)

Family ID: 96173918

Assignee: ADVANCED MICRO DEVICES, INC. (Santa Clara, CA)

Appl. No.: 18/397862

Filed: December 27, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250219948 A1	Jul. 03, 2025

Publication Classification

Int. Cl.: H04L47/20 (20220101); H04L47/215 (20220101); H04L47/32 (20220101)

U.S. Cl.:

CPC H04L47/20 (20130101); H04L47/215 (20130101); H04L47/32 (20130101);

Field of Classification Search

CPC: H04L (47/20); H04L (47/215); H04L (47/32)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9698791	12/2016	Vassiliev	N/A	N/A
9871733	12/2017	Atluri et al.	N/A	N/A
11516135	12/2021	Aibester et al.	N/A	N/A
2016/0248686	12/2015	Lee	N/A	H04W 4/24
2018/0234297	12/2017	Wang	N/A	H04L 41/0893
2022/0052936	12/2021	Chandrasekaran	N/A	H04L 41/40
2023/0064845	12/2022	Srinivasan	N/A	H04L 43/10
2023/0224217	12/2022	Subrahmanya	370/254	H04L 41/082
2024/0163230	12/2023	Danivas	N/A	H04L 49/3063

Primary Examiner: Thieu; Benjamin M

Background/Summary

TECHNICAL FIELD

(1) The systems and methods relate to computer networks, local area networks, networking devices that may be a router, a switch, a network interface card (NIC), a smartNIC, a distributed service card (DSC), etc. The systems and methods also relate to elements of networking devices such as semiconductor chips implementing packet processing pipeline circuits, programmable policer circuits, and aggregated token buckets.

BACKGROUND

(2) Networking devices process network flows by receiving network packets and processing the network packets. The network packets are often processed by examining the packet's header data and applying networking policies such as routing policies, firewall policies, load balancing policies, etc. Packet processing may be performed by a packet processing pipeline such as a “P4” packet processing pipeline. The concept of a domain-specific language for programming protocol-independent packet processors, known simply as “P4,” developed to provide some flexibility at the data plane of a networking device. The P4 domain-specific language for programming the data plane of networking devices has been defined in the “P4.sub.16 Language Specification,” published by the P4 Language Consortium. P4 (also referred to herein as the “P4 specification,” the “P4 language,” and the “P4 program”) may be implemented on a large variety of targets including switches, routers, programmable NICs, software switches, field programmable gate arrays (FPGAs), and application specific integrated circuits (ASICs). As described in the P4 specification,

the primary abstractions provided by the P4 language relate to header types, parsers, tables, actions, match-action units, control flow, extern objects, user-defined metadata, and intrinsic metadata. Packet processing pipeline circuits may be central processing unit (CPU) offloads that offload work from the CPUs and thereby free the CPUs to perform other tasks. Cryptographic transform circuits, compression circuits, and decompression circuits, are also examples of CPU offloads.

BRIEF SUMMARY OF SOME EXAMPLES

(3) The following presents a summary of one or more aspects of the present disclosure, to provide a basic understanding of such aspects. This summary is not an extensive overview of all contemplated features of the disclosure and is intended neither to identify key or critical elements of all aspects of the disclosure nor to delineate the scope of any or all aspects of the disclosure. Its sole purpose is to present some concepts of one or more aspects of the disclosure as a prelude to the more detailed description that is presented later.

(4) One aspect of the subject matter described in this disclosure may be implemented in a system. The system may include a programmable policer circuit configured to use a token bucket and a policing policy to produce a policing decision for a network packet in a network flow associated with the policing policy, a packet processing pipeline circuit configured to implement a data plane, and a processor configured to implement a control plane, wherein the programmable policer circuit, the packet processing pipeline circuit, and the processor are further configured to implement network flow policing that involves the control plane configuring the data plane to process the network packet, and the control plane associating the network flow with the policing policy such that the data plane is configured to process the network packet in accordance with the policing decision.

(5) Another aspect of the subject matter described in this disclosure may be implemented by a method. The method may include using a network packet to identify a programmable policer circuit and a policing policy for a network flow that includes the network packet, and obtaining a policing decision for the network packet from the programmable policer circuit, the policing decision in accordance with the policing policy. The method may also include processing, by a packet processing pipeline circuit, the network packet in accordance with the policing decision.

(6) Yet another aspect of the subject matter described in this disclosure may be implemented in a system. The system may include a programmable policing means for producing a policing decision for a network packet in accordance with a tokens value and a policing policy for a network flow, and a packet processing means for processing the network packet in accordance with the policing decision, wherein the programmable policing means produces a second policing decision for a second network packet in accordance with the tokens value and a second policing policy for a second network flow, and the packet processing means processes the second network packet in accordance with the second policing decision.

(7) In some implementations of the methods and devices, the packet processing pipeline circuit is configured to use the network packet to identify a flow table entry that includes a policing policy identifier identifying the policing policy and a policer identifier identifying the programmable policer circuit, and the packet processing pipeline circuit uses the policer identifier to provide policing data to the programmable policer circuit, the policing data including the policing policy identifier. In some implementations of the methods and devices, the control plane uses a plurality of networking policies to produce a plurality of flow table entries that includes the flow table entry. In some implementations of the methods and devices, the token bucket is an aggregated token bucket, the programmable policer circuit is configured to use the aggregated token bucket and a second policing policy to produce a second policing decision for a second network packet in a second network flow associated with the second policing policy, and the packet processing pipeline circuit is configured to process the second network packet in accordance with the second policing decision. In some implementations of the methods and devices, the network flow is directed to or from a network endpoint and the second network flow is directed to or from the network endpoint.

In some implementations of the methods and devices, wherein the network endpoint is a virtual machine.

(8) In some implementations of the methods and devices, the packet processing pipeline circuit includes a parser stage and a plurality of match-action processing stages, the parser stage produces a packet header vector (PHV) from the network packet, and one of the match-action processing stages uses the PHV to identify the flow table entry. In some implementations of the methods and devices, the policing decision is to drop the network packet in response to determining that allowing the network packet would cause a tokens value of the aggregated token bucket to be less than a threshold value, and the second policing decision is to allow the second network packet even though the tokens value is below the threshold value. In some implementations of the methods and devices, the network flow is a non-storage flow, and the second network flow is a storage flow. In some implementations of the methods and devices, the policing decision is to drop the network packet in response to determining that allowing the network packet would cause a tokens value of the aggregated token bucket to be less than a threshold value, and the second policing policy is to always allow and to always decrease the tokens value based on a size of an allowed network packet.

(9) In some implementations of the methods and devices, a flow table includes a flow table entry that is associated with first network flow, the packet processing pipeline circuit uses the network packet to identify the flow table entry, the flow table entry includes a policing policy identifier that identifies the policing policy and includes a policer identifier that identifies the programmable policer circuit, and the packet processing pipeline circuit uses the policer identifier and the policing policy identifier to obtain the policing decision from the programmable policer circuit. In some implementations of the methods and devices, the method further includes using a plurality of networking policies to produce a plurality of flow table entries that includes the flow table entry. In some implementations of the methods and devices, the method further includes using a tokens value in an aggregated token bucket for policing the network flow and a second network flow that is associated with a second policing policy, using, by the programmable policer circuit, the aggregated token bucket and the second policing policy to produce a second policing decision for a second network packet that is in the second network flow, and processing, by the packet processing pipeline circuit, the second network packet in accordance with the second policing decision.

(10) In some implementations of the methods and devices, the system further includes a lookup means for using the network packet to identify the programmable policing means and the policing policy, wherein the lookup means uses the second network packet to identify the programmable policing means and the second policing policy.

(11) These and other aspects will become more fully understood upon a review of the detailed description, which follows. Other aspects and features will become apparent to those of ordinary skill in the art, upon reviewing the following description of specific examples in conjunction with the accompanying figures. While features may be discussed relative to certain examples and figures below, any example may include one or more of the advantageous features discussed herein. In other words, while one or more examples may be discussed as having certain advantageous features, one or more of such features may also be used in accordance with the various examples discussed herein. In similar fashion, while the examples may be discussed below as devices, systems, or methods, the examples may be implemented in various devices, systems, and methods.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a high-level conceptual diagram illustrating a system using programmable policer circuits to implement strict priority for a class of traffic, according to some aspects.

- (2) FIG. 2 is a functional block diagram of a networking device having a control plane and a data plane and in which aspects may be implemented.
- (3) FIG. 3 is a functional block diagram illustrating an example of a match-action processing stage in a match-action pipeline according to some aspects.
- (4) FIG. 4 is a functional block diagram of a networking device having a semiconductor chip such as an application specific integrated circuit (ASIC) or field programmable gate array (FPGA), according to some aspects.
- (5) FIG. 5 is a high-level diagram illustrating an example of generating a packet header vector from a packet according to some aspects.
- (6) FIG. 6 illustrates a block diagram of a match processing unit (MPU) that may be used within the exemplary system of FIG. 4 to implement some aspects.
- (7) FIG. 7 illustrates a block diagram of a packet processing pipeline circuit that may be included in the exemplary system of FIG. 4.
- (8) FIG. 8 illustrates packet headers and payloads of packets for network flows including a NVMe/TCP PDU in a layer 4 payload according to some aspects.
- (9) FIG. 9 illustrates a high-level flow diagram of a packet processing pipeline obtaining a policing decision from a programmable policer circuit, according to some aspects.
- (10) FIG. 10 is a high-level conceptual diagram illustrating a metering circuit refilling the tokens in a token bucket that may be used as an aggregated token bucket, according to some aspects.
- (11) FIG. 11 is a high-level conceptual diagram illustrating policing policy identifiers that are associated with policing policies, according to some aspects.
- (12) FIG. 12 is a high-level flow diagram illustrating an example of a process for producing policing decisions and that may be implemented by a programmable policer circuit, according to some aspects.
- (13) FIG. 13 is a high-level conceptual diagram illustrating an example of allowing a packet received from a network endpoint, according to some aspects.
- (14) FIG. 14 is a high-level conceptual diagram illustrating an example of dropping a packet received from a network endpoint, according to some aspects.
- (15) FIG. 15 is a high-level conceptual diagram illustrating a control plane using networking policies to produce flow table entries, according to some aspects.
- (16) FIG. 16 is a high-level flow diagram illustrating a method for using programmable policer circuits to implement strict priority for a class of traffic, according to some aspects.
- (17) FIG. 17 is another high-level flow diagram illustrating a method for using programmable policer circuits to implement strict priority for a class of traffic, according to some aspects.
- (18) FIG. 18 is a high-level block diagram illustrating an example of a programmable policer circuit, according to some aspects.
- (19) Throughout the description, similar reference numbers may be used to identify similar elements.

DETAILED DESCRIPTION

(20) It will be readily understood that the components of the examples as generally described herein and illustrated in the appended figures could be arranged and designed in a wide variety of different configurations. Thus, the following more detailed description of various examples, as represented in the figures, is not intended to limit the scope of the present disclosure but is merely representative of various examples. While the various aspects of the examples are presented in drawings, the drawings are not necessarily drawn to scale unless specifically indicated.

(21) Systems and methods that implement aspects may have various differing forms. The described systems and methods are to be considered in all respects only as illustrative and not restrictive. The scope of the claims is, therefore, indicated by the claims themselves rather than by this detailed description. All changes which come within the meaning and range of equivalency of the claims are to be embraced within their scope.

(22) Reference throughout this specification to features, advantages, or similar language does not imply that any system or method implements each and every aspect that may be realized. Rather, language referring to the features and advantages is understood to mean that a specific feature, advantage, or characteristic described in an example may be implemented in or by at least one example. Thus, discussions of the features and advantages, and similar language, throughout this specification may, but do not necessarily, refer to the same example.

(23) Furthermore, the described features, advantages, characteristics, and aspects may be combined in any suitable manner in one or more systems or methods. One skilled in the relevant art will recognize, in light of the description herein, that one example may be practiced without one or more of the specific features or advantages of another example. In other instances, additional features and advantages may be recognized in one example that may not be present in all the examples.

(24) Reference throughout this specification to “one example”, “an example”, or similar language means that a particular feature, structure, or characteristic described in connection with the indicated example is included in at least one example. Thus, the phrases “in one example”, “in an example”, and similar language throughout this specification may, but do not necessarily, all refer to the same example.

(25) In the field of data networking, the functionality of networking devices such as switches, routers, and NICs are often described in terms of functionality that may be associated with a “control plane” and functionality that may be associated with a “data plane.” In general, the control plane refers to components and/or operations that are involved in managing forwarding information and the data plane refers to components and/or operations that are involved in forwarding packets from an input interface to an output interface according to the forwarding information provided by the control plane. The data plane may also refer to components and/or operations that implement packet processing operations related to encryption, decryption, compression, decompression, firewalling, and telemetry.

(26) Aspects described herein may use a match-action pipeline in the data plane to process packets. A match-action pipeline may process network flows extremely quickly when configured to process those network flows. A match-action pipeline has a series of match-action stages. Upon receiving a packet of a network flow, the first match-action stage of a match-action pipeline may generate an index from data in the packet header. Finding a flow table entry for the network flow at the index location in the flow table is the “match” portion of “match-action”. If there is a “match”, the “action” is performed to thereby process the packet. If there is no flow table entry for the network flow, it is a new network flow that the match-action pipeline is not yet configured to process. If there is no match, then the match-action stage may perform a default action. Match-action stages may pass the packet to a subsequent match-action stage in the match-action pipeline.

(27) The high-volume and rapid decision-making that occurs at the data plane is often implemented in one or more semiconductor chips. An application specific integrated circuit (ASIC) and a field programmable gate array (FPGA) are examples of such semiconductor chips. A semiconductor chip such as a fixed function ASIC may enable high-volume and rapid packet processing. However, a fixed function ASIC typically does not provide enough flexibility to adapt to changing needs. Data plane processing implemented in FPGAs may provide a high level of flexibility in data plane processing.

(28) Data centers run workloads for a variety of tenants. The workloads may include one or more virtual machine (VM) running on a host computer provisioned by the data center. An aspect of hosting those workloads is meeting quality of service (QoS) guarantees. Another aspect is limiting tenants to the level of resources they have contracted for. For example, a first tenant that has contracted for 1 GB/s of bandwidth may attempt to consume 10 GB/s of bandwidth. A second tenant that has contracted for 10 GB/s of bandwidth may get less than that bandwidth when the first tenant is allowed to consume too much of the available bandwidth. As such, communications for

some workloads are throttled so that others may receive their allotment of network services. The communications for a workload may be throttled by a NIC installed in the host running the workload. The NIC may include policer circuits that may track the resource consumption (e.g., network bandwidth, packets per second, etc.) of the network flows and may produce a policing decision to drop a network packet in response to determining that a network flow will exceed its allocation of resources. Different types of network flows (e.g., storage flows and non-storage flows) may be policed differently. A storage flow is a network flow between two network endpoints where one of the endpoints is providing storage (e.g., network attached storage devices, storage area networks, etc.). Network flows that are not storage flows are non-storage flows. For example, a network flow between a web server and a user device may be a non-storage flow. A policer circuit, unlike the programmable policer circuits discussed herein, implements a single policing policy such as “drop all packets that would result in a threshold usage being exceeded”.

(29) In some cases, the number of policer circuits may limit the number of network flows that may be concurrently policed or the total network resource consumption of a tenant or workload may be policed instead of policing the flows individually. A policer circuit may implement a single policing policy whereas different policing policies based on aggregated network resource consumption may be needed for different types of network flows.

(30) In some examples, a tenant's or workload's aggregated network resource consumption may be policed with a programmable policer circuit, a token bucket, and a packet processing pipeline circuit. An aspect of the programmable policer circuit is that it may implement different policing policies for different network flows. A token bucket shared by multiple network flows may be referred to as an aggregated token bucket and may be used for policing a tenant's or workload's aggregated network resource consumption. A networking device (e.g., a SmartNIC) may include a data plane implemented by a packet processing pipeline circuit (e.g., a P4 pipeline circuit) and a control plane implemented by one or more general purpose processors (e.g., an x86 or ARM CPUs). The control plane may configure the data plane to process the network packets of network flows and may associate different network flows with different policing policies. The control plane may also associate one of the networking device's programmable policer circuits with a group of network flows that have different policing policies. An aggregated token bucket may be used to track the resource consumption of all the network flows associated with one programmable policer circuit. The programmable policer circuit may check the aggregated token bucket and produce a policing decision for a network packet by implementing the policing policy associated with the network flow that includes the network packet. The packet processing pipeline circuit may act on the policing decisions by dropping or not dropping network packets.

(31) One advantage over a non-programmable policer circuit is that a programmable policer circuit may, while using a single aggregated token bucket, produce different policing decisions for network packets in different network flows. Another advantage is that a programmable policer circuit may implement strict priority for a class of traffic (e.g., storage flows) by allowing that class of traffic when a different class of traffic (e.g., non-storage flows) is dropped. Yet another advantage is that a single programmable policer circuit may be used for a workload's different classes of network traffic, thereby conserving the number of policer circuits that are used. Still yet another advantage is that different network flows may be policed using different policing policies, which may improve the performance of the workloads. For example, the execution of a workload may be paused while the workload is waiting for a storage packet but the workload's execution may be more resilient with respect to non-storage network flows. A more specific example is a workload such as a web server that does not wait for responses to outgoing packets HTTP packets but pauses while waiting for a response to a read request sent to a storage area network (SAN).

(32) FIG. 1 is a high-level conceptual diagram illustrating a system using programmable policer circuits to implement strict priority for a class of traffic, according to some aspects. Two network packets are received. The first network packet **101** is in a first network flow. The second network

packet **102** is in a second network flow. A packet processing pipeline circuit **103** may process the network packets. A parser **104** produces a packet header vector (PHV) **105** from a network packet. A first processing stage **106**, which may also be referred to as the first match-action stage, may use the PHV to look up the flow table entry for a network packet. The flow table **110** includes a first flow table entry **111** for the first network flow and includes a second flow table entry **114** for the second network flow. As such, the first processing stage uses the first flow table entry **111** for processing the first network packet **101** and uses the second flow table entry **114** for processing the second network packet **102**. A flow table entry may include a field for a policer circuit identifier and a field for a policing policy identifier. The policer circuit identifier **112** of the first flow table entry **111** identifies a programmable policer circuit **123** that is one of the programmable policer circuits. The policer circuit identifier **115** of the second flow table entry **114** also identifies the programmable policer circuit **123**. As such, the same programmable policer circuit **123** is used for the first network packet **101** and for the second network packet **102** even though the network packets are in different network flows. The packet processing pipeline circuit **103** and the programmable policer circuit may be implemented in a semiconductor chip such as the semiconductor chip **401** illustrated in FIG. 4.

(33) The policing policy identifier **113** of the first flow table entry **111** identifies an “accumulate, don't drop” policing policy. The policing policy identifier **116** of the second flow table entry **114** identifies a “don't accumulate, drop” policing policy. The policing policy identifier **116** of the second flow table entry **114** also identifies the programmable policer circuit **123**. As such, the programmable policer circuit **123** will implement one policing policy for the first network packet **101** and will implement a different policing policy for the second network packet **102**. The first processing stage **106** may add the policing policy identifier and the policer circuit identifier to the PHV **107**, which is passed to the next stage of the packet processing pipeline circuit **103**. The second processing stage **108**, which may also be called the second match-action stage, receives the PHV and sends policing data **120** to the programmable policer circuit identified by the policer circuit identifier in the flow table entry. The policing data **120** may include the packet size **121** and the policing policy identifier **122** from the flow table entry. The packet size **121** is a resources required value. A resources required value indicates the amount of resources that would be required for processing a network packet. The programmable policer circuit **123** uses an aggregated token bucket **140** to track the amount of network resources consumed. For example, packet size **121** may be used to ensure that bandwidth limits are respected. The tokens value **141** in the aggregated token bucket indicates an amount of network resources that are available for the network flows. The programmable policer circuit **123** may use the tokens value **141** and the packet size **121** to produce a policing decision **124**. The second processing stage **108** may implement the policing decision by dropping the network packet if the policing decision **124** is “drop” or by continuing to process the network packet if the policing decision **124** is “allow”. When the policing decision for a network packet is “allow”, the network packet is an allowed network packet. A processed network packet **130** may be the result when the packet processing pipeline continues processing the network packet. In some implementations, the first processing stage **106** may read the flow table entry, send the policing data to the programmable policer circuit, and implement the policing decision.

(34) The aggregated token bucket **140** illustrated in FIG. 1 may have a threshold value that equals 0. The programmable policer circuit may compare the difference of the tokens value and the packet size (e.g., tokens value minus packet size) to the threshold value and make the policing decision based on the result of the comparison. The threshold value may be implicit. For example, the threshold may be assumed to be zero. Alternatively, a threshold value may be associated with or stored in the aggregated token bucket or the programmable policer circuit. The flow table entries may include a threshold value that may then be included in the policing data and then used by the policer circuit to make the policing decision.

(35) FIG. 2 is a functional block diagram of a networking device having a control plane and a data

plane and in which aspects may be implemented. A networking device **201** may have a control plane **203** and a data plane **202**. The control plane may be implemented by a general purpose processor such as a central processing unit (CPU) containing one or more CPU cores. The control plane provides forwarding information (e.g., in the form of table management information or configuration data) to the data plane and the data plane receives packets on input interfaces, processes the received packets, and then forwards packets to desired output interfaces. Additionally, control traffic (e.g., in the form of packets) may be communicated from the data plane to the control plane and/or from the control plane to the data plane. The data plane and control plane are sometimes referred to as the “fast” plane and the “slow” plane, respectively. In general, the control plane is responsible for less frequent and less time-sensitive operations such as updating Forwarding Information Bases (FIBs) and Label Forwarding Information Bases (LFIBs), while the data plane is responsible for a high volume of time-sensitive forwarding decisions that need to be made at a rapid pace. The control plane may implement operations related to packet routing that include InfiniBand channel adapter management functions, Open Shortest Path First (OSPF), Enhanced Interior Gateway Routing Protocol (EIGRP), Border Gateway Protocol (BGP), Intermediate System to Intermediate System (IS-IS), Label Distribution Protocol (LDP), routing tables and/or operations related to packet switching that include Address Resolution Protocol (ARP) and Spanning Tree Protocol (STP). The data plane (which may also be referred to as the “forwarding” plane) may implement operations related to parsing packet headers, Quality of Service (QoS), filtering, encapsulation, queuing, and policing. Although some functions of the control plane and data plane are described, other functions may be implemented in the control plane and/or the data plane.

(36) Some techniques exist for providing flexibility at the data plane of networking devices that are used in data networks. For example, the concept of a domain-specific language for programming protocol-independent packet processors, known simply as “P4,” has developed to provide some flexibility at the data plane of a networking device. The document “P4.sub.16 Language Specification,” published by the P4 Language Consortium is well known in the field and describes the P4 domain-specific language that may be used for programming the data plane of networking devices. P4 (also referred to herein as the “P4 specification,” the “P4 language,” and the “P4 program”) is designed to be implementable on a large variety of targets including switches, routers, programmable NICs, software switches, semiconductor chip, FPGAs, and ASICs. As described in the P4 specification, the primary abstractions provided by the P4 language relate to header types, parsers, tables, actions, match-action units, match-action pipeline stages, control flow, extern objects, user-defined metadata, and intrinsic metadata.

(37) The data plane **202** includes multiple receive (RX) media access controllers (MACs) **211** and multiple transmit (TX) MACs **210**. The RX MACs **211** implement media access control on incoming packets via, for example, a layer 2 protocol such as Ethernet. The layer 2 protocol may be Ethernet and the RX MACs may be configured to implement operations related to, for example, receiving frames, half-duplex retransmission and back-off functions, Frame Check Sequence (FCS), interframe gap enforcement, discarding malformed frames, and removing the preamble, Start Frame Delimiter (SFD), and padding from a packet. Likewise, the TX MACs **210** implement media access control on outgoing packets via, for example, Ethernet. The TX MACs may be configured to implement operations related to, for example, transmitting frames, half-duplex retransmission, and back-off functions, appending an FCS, interframe gap enforcement, and prepending a preamble, an SFD, and padding.

(38) As illustrated in FIG. 2, a P4 program is provided to the data plane **202** via the control plane **203**. Communications between the control plane and the data plane may use a dedicated channel or bus, may use shared memory, etc. The P4 program includes software code that configures the functionality of the data plane **202** to implement particular processing and/or forwarding logic and to implement processing and/or forwarding tables that are populated and managed via P4 table

management information that is provided to the data plane from the control plane. Control traffic (e.g., in the form of packets) may be communicated from the data plane to the control plane and/or from the control plane to the data plane. In the context of P4, the control plane corresponds to a class of algorithms and the corresponding input and output data that are concerned with the provisioning and configuration of the data plane. The data plane corresponds to a class of algorithms that describe transformations on packets by packet processing systems.

(39) The data plane **202** includes a programmable packet processing pipeline **204** that is programmable using a domain-specific language such as P4. As described in the P4 specification, a programmable packet processing pipeline may include an arbiter **205**, a parser **206**, a match-action pipeline **207**, a deparser **208**, and a demux/queue **209**. The data plane elements described may be implemented as a P4 programmable switch architecture, as a P4 programmable NIC, as a P4 programmable router, or some other architecture. The arbiter **205** may act as an ingress unit receiving packets from RX MACs **211** and may also receive packets from the control plane via a control plane packet input **212**. The arbiter **205** may also receive packets that are recirculated to it by the demux/queue **209**. The demux/queue **209** may act as an egress unit and may also be configured to send packets to a drop port (the packets thereby disappear), to the arbiter via recirculation, and to the control plane **203** via an output central processing unit (CPU) port **213**. The arbiter **205** and the demux/queue **209** may be configured through the domain-specific language (e.g., P4).

(40) The parser **206** is a programmable element that may be configured through the domain-specific language (e.g., P4) to extract information from a packet (e.g., information from the header of the packet). As described in the P4 specification, parsers describe the permitted sequences of headers within received packets, how to identify those header sequences, and the headers and fields to extract from packets. The information extracted from a packet by the parser may be referred to as a packet header vector (PHV). The parser may identify certain fields of the header and may extract the data corresponding to the identified fields to generate the PHV. The PHV may include other data (often referred to as “metadata”) that is related to the packet but not extracted directly from the header, including for example, the port or interface on which the packet arrived at the networking device. Thus, the PHV may include other packet related data (metadata) such as input/output port number, input/output interface, or other data in addition to information extracted directly from the packet header. The PHV produced by the parser may have any size or length. For example, the PHV may be at least 4 bits, 8 bits, 16 bits, 32 bits, 64 bits, 128 bits, 256 bits, or 512 bits. In some cases, a PHV having even more bits (e.g., 6 Kb) may include all relevant header fields and metadata corresponding to a received packet. The size or length of a PHV corresponding to a packet may vary as the packet passes through the match-action pipeline.

(41) The deparser **208** is a programmable element that is configured through the domain-specific language (e.g., P4) to generate packet headers from PHVs at the output of match-action pipeline **207** and to construct outgoing packets by reassembling the header(s) such as Ethernet headers, internet protocol (IP) headers, InfiniBand protocol data units (PDUs), etc. as determined by the match-action pipeline. In some cases, a packet/payload may travel in a separate queue or buffer **220**, such as a first-in-first-out (FIFO) queue, until the packet payload is reassembled with its corresponding PHV at the deparser to form a packet. The deparser may rewrite the original packet according to the PHV fields that have been modified (e.g., added, removed, or updated). In some cases, a packet processed by the parser may be placed in a packet buffer/traffic manager for scheduling and possible replication. In some cases, once a packet is scheduled and leaves the packet buffer/traffic manager, the packet may be parsed again to generate an egress PHV. The egress PHV may be passed through a match-action pipeline after which a final deparser operation may be executed (e.g., at deparser **208**) before the demux/queue **209** sends the packet to the TX MAC **210** or recirculates it back to the arbiter **205** for additional processing.

(42) A networking device **201** may have a peripheral component interconnect extended (PCIe)

interface such as PCIe media access control (MAC) **214**. A PCIe MAC may have a base address register (BAR) at a base address in a host system's memory space. Processes, typically device drivers within the host system's operating system, may communicate with a NIC via a set of registers beginning with the BAR. Some PCIe devices are single root input output virtualization (SR-IOV) capable. Such PCIe devices may have a physical function (PF) and a virtual function (VF). A PCIe SR-IOV capable device may have multiple VFs. A PF BAR map **215** may be used by the host machine to communicate with the PCIe card. A VF BAR map **216** may be used by a virtual machine (VM) running on the host to communicate with the PCIe card. Typically, the VM may access the NIC using a device driver within the VM and at a memory address within the VMs memory space. Many SR-IOV capable PCIe cards may map that location in the VM's memory space to a VF BAR. As such a VM may be configured as if it has its own NIC while in reality it is associated with a VF provided by a SR-IOV capable NIC. As discussed below, some PCIe devices may have multiple PFs. For example, a NIC may provide network connectivity via one PF and may provide an InfiniBand channel adapter via another PF. As such, the NIC may provide "NIC" VFs and "InfiniBand" VFs to VMs running on the host. The InfiniBand PF and VFs may be used for data transfers, such as remote direct memory access (RDMA) transfers to other VMs running on the same or other host computers. Similarly, a NIC may provide non-volatile memory express (NVMe) and small computer system interface (SCSI) PFs and VFs to VMs running on the host.

(43) FIG. **3** is a functional block diagram illustrating an example of a match-action unit **301** in a match-action pipeline **300** according to some aspects. FIG. **3** introduces certain concepts related to match-action units and match-action pipelines and is not intended to be limiting. The match-action units are processing stages, often called stages or match-action processing stages, of the packet processing pipeline. The match-action processing stages **301**, **302**, **303** of the match-action pipeline **300** are programmed to perform "match-action" operations in which a match unit performs a lookup using at least a portion of the PHV and an action unit performs an action based on an output from the match unit. A PHV generated at the parser may be passed through each of the match-action processing stages in the match-action pipeline in series and each match-action processing stage may implement a match-action operation or policy. The PHV and/or table entries may be updated in each stage of match-action processing according to the actions specified by the P4 programming. In some instances, a packet may be recirculated through the match-action pipeline, or a portion thereof, for additional processing. The first match-action processing stage **301** receives the first PHV **305** as an input and outputs the second PHV **306**. The second match-action processing stage **302** receives the second PHV **306** as an input and outputs the third PHV **307**. The third match-action processing stage **303** receives the third PHV **307** as an input and outputs the fourth PHV **308**. The match-action processing stages are arranged as a match-action pipeline that passes the PHVs from one match-action processing stage to the next match-action processing stage in the pipeline.

(44) An expanded view of elements of a match-action processing stage **301** of match-action pipeline **300** is shown. The match-action processing stage includes a match unit **317** (also referred to as a "table engine") that operates on an input PHV **305** and an action unit **314** that produces an output PHV **306**, which may be a modified version of the input PHV **305**. The match unit **317** may include key construction logic **309**, a lookup table **310**, and selector logic **312**. The key construction logic **309** is configured to generate a key from at least one field in the PHV (e.g., 5-tuple, InfiniBand queue pair identifiers, etc.). The lookup table **310** is populated with key-action pairs, where a key-action pair may include a key (e.g., a lookup key) and corresponding action code **315** and/or action data **316**. A P4 lookup table may be viewed as a generalization of traditional switch tables, and may be programmed to implement, for example, routing tables, flow lookup tables, access control lists (ACLs), and other user-defined table types, including complex multi-variable tables. The key generation and lookup functions constitute the "match" portion of the operation and produce an action that is provided to the action unit via the selector logic. The action

unit executes an action over the input data (which may include data **313** from the PHV) and provides an output that forms at least a portion of the output PHV. For example, the action unit executes action code **315** on action data **316** and data **313** to produce an output that is included in the output PHV **306**. If no match is found in the lookup table, then a default action **311** may be implemented. A flow miss is an example of a default action that may be executed when no match is found. The operations of the match-action processing stages may be programmable by the control plane via P4 and the contents of the lookup table (e.g., a flow table) may be managed by the control plane.

(45) FIG. 4 is a functional block diagram of a networking device **430** having a semiconductor chip **401** such as an ASIC or FPGA, according to some aspects. The semiconductor chip **401** shows a single semiconductor chip implementing a large number of hardware functions. A different and substantially equivalent implementation may employ a chiplet architecture. If the networking device is a network interface card (NIC) then the NIC may be installed in a host computer and may act as a networking device for the host computer and for virtual machines running on the host computer. Such a NIC may have a PCIe connection **431** for communicating with the host computer via a host PCIe connection. The networking device **430** may have a semiconductor chip **401**, off-chip memory **432**, and ethernet ports **433**. The off-chip memory **432** may be one of the widely available memory modules or chips such as double data rate 5 (DDR5) synchronous dynamic random-access memory (SDRAM) such that the semiconductor chip **401** has access to many gigabytes of memory on the networking device **430**. The ethernet ports **433** provide physical connectivity to a computer network such as the internet. The NIC may include a printed circuit board to which the semiconductor chip **401** and the memory **432** are attached.

(46) The semiconductor chip may have many core circuits interconnected by an on-chip communications fabric, sometimes called a network on a chip (NOC) **402**. NOCs are often implementations of standardized communications fabrics such as the widely used advanced extensible interface (AXI) bus. The semiconductor chip's core circuits may include a PCIe interface **427**, CPU **403**, first packet processing pipeline circuit **408**, memory interface circuit **415**, on-chip memory **416** that may be a static random access memory (SRAM), programmable policer circuits **417**, a packet buffer **422**, metering circuit **423**, and packet ingress/egress circuits **414**. The PCIe interface **427** may be used to communicate with a host computer via the PCIe connection **431**. The CPU **403** may include numerous CPU cores such as a first CPU core **405**, a second CPU core **406**, and a third CPU core **407**. The CPU **403** may implement the control plane of the networking device **430**. The first packet processing pipeline circuit **408** may include a pipeline ingress circuit **413**, a parser circuit **412**, match-action pipeline **411**, a deparser circuit **410**, and a pipeline egress circuit **409**. The second packet processing pipeline circuit **425** may include a PHV ingress circuit **428**, a match-action pipeline **434**, and a direct memory access (DMA) output circuit **426**. The programmable policer circuits **417** may include a first programmable policer circuit **418**, a second programmable policer circuit **419**, and a last programmable policer circuit. A programmable policer circuit is a hardware implemented circuit that may be used for policing network flows. The specific core circuits implemented within the non-limiting example of the semiconductor chip **401** may be selected such that the semiconductor chip implements many, perhaps all, of the functionality of an InfiniBand channel adapter, of an NVMe card, and of a networking device that processes network flows carried by internet protocol (IP) packets.

(47) A network device may include precision clocks that output a precise time, clocks that are synchronized to remote authoritative clocks via precision time protocol (PTP), and hardware clocks **424**. A hardware clock may provide a time value (e.g., year/day/hour/minute/second/ . . .) or may simply be a counter that is incremented by one at regular intervals (e.g., once per clock cycle for a device having a 10 nsec. clock period). Time values obtained from the clocks may be used as timestamps for events such as enqueueing/dequeueing a packet.

(48) The first packet processing pipeline circuit **408** is a specialized set of elements for processing

PHVs including PHVs for network packets such as internet protocol (IP) packets and InfiniBand protocol data units (PDUs). The data plane may be implemented by a packet processing pipeline circuit such as the first packet processing pipeline circuit **408**. The first packet processing pipeline circuit **408** may be a P4 packet processing pipeline circuit that implements a P4 pipeline that may be configured using a domain-specific language such as the P4 domain specific language. As described in the P4 specification, the primary abstractions provided by the P4 language relate to header types, parsers, tables, actions, match-action units, control flow, extern objects, user-defined metadata, and intrinsic metadata.

(49) The second packet processing pipeline circuit **425** is a specialized set of elements for processing PHVs including PHVs for network packets such as internet protocol (IP) packets and InfiniBand protocol data units (PDUs). The second packet processing pipeline circuit **425** may be a P4 packet processing pipeline circuit that implements a P4 pipeline that may be configured using a domain-specific language such as the P4 domain specific language. As described in the P4 specification, the primary abstractions provided by the P4 language relate to header types, parsers, tables, actions, match-action units, control flow, extern objects, user-defined metadata, and intrinsic metadata. The data plane may be implemented by the first packet processing pipeline circuit **408** in combination with the second packet processing pipeline circuit **425**.

(50) The networking device **430** may include a memory **432** for running Linux or some other operating system and for storing data used by the processes implementing network services, upgrading the control plane, and upgrading the data plane. The networking device may use the memory **432** to store a flow table **440**. The flow table may include flow table entries such as a first flow table entry **441**, a second flow table entry **442**, and a last flow table entry. The memory **432** may also store aggregated token buckets **445** such as a first aggregated token bucket **446**, a second aggregated token bucket **447**, and a last aggregated token bucket. In some implementations, flow tables or aggregated token buckets are stored in the on-chip memory **416** of the semiconductor chip **401** such that the token buckets and flow table entries may be more rapidly accessed by the packet processing pipeline circuits **408**, **425** and the metering circuit **423** illustrated in FIG. **10**.

(51) The CPU cores **405**, **406**, **407** may be general purpose processor cores, such as ARM processor cores, microprocessor without interlocked pipelined stages (MIPS) processor cores, and/or x86 processor cores, as is known in the field. Each CPU core may include an arithmetic logic unit (ALU), a register bank, an instruction fetch unit, and an instruction decoder, which are configured to execute instructions independently of the other CPU cores. The CPU cores may be Reduced Instruction Set Computers (RISC) CPU cores that are programmable using a general-purpose programming language such as C.

(52) The CPU **403** may also include a bus interface, internal memory, and a memory management unit (MMU) and/or memory protection unit. For example, the CPU cores may include internal cache, e.g., L1 cache and/or L2 cache, and/or may have access to nearby L2 and/or L3 cache. Each CPU core may include core-specific L1 cache, including instruction-cache and data-cache and L2 cache that is specific to each CPU core or shared amongst a small number of CPU cores. L3 cache may also be available to the CPU cores.

(53) There may be multiple CPU cores **405**, **406**, **407** available for control plane functions and for implementing aspects of a slow data path that includes software implemented packet processing functions. The CPU cores may be used to implement discrete packet processing operations such as L7 applications (e.g., HTTP load balancing, L7 firewalling, and/or L7 telemetry), certain InfiniBand channel adapter functions, flow table insertion or table management events, connection setup/management, multicast group join, deep packet inspection (DPI) (e.g., URL inspection), storage volume management (e.g., NVMe volume setup and/or management), encryption, decryption, compression, and decompression, which may not be readily implementable through a domain-specific language such as P4, in a manner that provides fast path performance as is expected of data plane processing.

(54) The packet buffer **422** may act as a central on-chip packet switch that delivers packets from the network interfaces **433** to packet processing elements of the data plane and vice-versa. The packet processing elements may include a slow data path implemented in software and a fast data path implemented by a packet processing pipeline circuits **408**, **425**.

(55) The first packet processing pipeline circuit **408** may be a specialized circuit or part of a specialized circuit using one or more semiconductor chips such as ASICs or FPGAs to implement programmable packet processing pipelines such as the programmable packet processing pipeline **204** of FIG. 2. Some networking devices include semiconductor chips such as ASICs or FPGAs implementing a P4 pipeline of a data plane within the networking device.

(56) All data transactions in the semiconductor chip **401**, including on-chip memory transactions, and register reads/writes may be performed via a coherent interconnect **402**. In one non-limiting example, the coherent interconnect may be provided by a network on a chip (NOC) "IP core". Semiconductor chip designers may license and use prequalified IP cores within their designs. Prequalified IP cores may be available from third parties for inclusion in chips produced using certain semiconductor fabrication processes. A number of vendors provide NOC IP cores. The NOC may provide cache coherent interconnect between the NOC masters, including the first packet processing pipeline circuit **408**, the second packet processing pipeline circuit **425**, CPU **403**, memory interface circuit **415**, and PCIe interface **427**. The interconnect may distribute memory transactions across a plurality of memory interfaces using a programmable hash algorithm. All traffic targeting the memory may be stored in a NOC cache (e.g., 1 MB cache). The NOC cache may be kept coherent with the CPU core caches.

(57) FIG. 5 is a high-level diagram illustrating an example of generating an ingress packet header vector **506** from a packet **501** according to some aspects. The PHV **506** is an ingress PHV that is produced by a parser **502** parsing a packet **501** received via an ingress port as a bit stream. The parser **502** may receive a packet **501** that has layer 2, layer 3, layer 4, and layer 7 headers and payloads. The parser may generate a packet header vector (PHV) from packet **501**. The packet header vector **506** may include many data fields including data from packet headers **507** and metadata **522**. The metadata **522** may include data generated by the networking device such as the hardware port on which the packet **501** was received and the packet timestamps indicating when the packet **501** was received by the networking device, enqueued, dequeued, etc. The metadata **522** may also include data produced by the networking device while processing a packet or assembling a packet. Such metadata **522** may include a PHV type **525** (e.g., Ingress PHV).

(58) The source MAC address **508** and the destination MAC address **509** may be obtained from the packet's layer 2 header. The source IP address **511** may be obtained from the packet's layer 3 header. The source port **512** may be obtained from the packet's layer 4 header. The protocol **513** may be obtained from the packet's layer 3 header. The destination IP address **514** may be obtained from the packet's layer 3 header. The destination port **515** may be obtained from the packet's layer 4 header. The packet quality of service parameters **516** may be obtained from the packet's layer 3 header or another header based on implementation specific details. The layer 4 header data **517** may be obtained from the packet's layer 4 header. The multi-protocol label switching (MPLS) data **518**, such as an MPLS label, may be obtained from the packet's layer 2 header. The layer 7 header data **519** may be obtained from the packet's layer 7 header. The other layer 7 data fields **520** may be obtained from the packet's layer 7 payload. The other header information **521** is the other information contained in the packet's layer 2, layer 3, layer 4, and layer 7 headers.

(59) The packet 5-tuple **510** is often used for generating keys for looking up and reading entries in key-value tables such as flow tables. The packet 5-tuple **510** may include the source IP address **511**, the source port **512**, the protocol **513**, the destination IP address **514**, and the destination port **515**.

(60) Those practiced in computer networking protocols realize that the headers carry much more information than that described here, realize that substantially all of the headers are standardized by documents detailing header contents and fields, and know how to obtain those documents. The

parser may also be configured to output a payload **505**. Recalling that the parser **502** is a programmable element that is configured through the domain-specific language (e.g., P4) to extract information from a packet, the specific contents of the payload **505** are those contents specified via the domain specific language. For example, the contents of the payload **505** may be the layer 4 payload, the layer 7 payload, etc.

(61) FIG. **6** illustrates a block diagram of a match processing unit (MPU) **601**, also referred to as an action unit, that may be used within the exemplary system of FIG. **4** to implement some aspects. The MPU **601** may have multiple functional units, memories, and a register file. For example, the MPU **601** may have an instruction fetch unit **605**, a register file unit **606**, a communication interface **602**, arithmetic logic units (ALUs) **607** and various other functional units.

(62) In the illustrated example, the MPU **601** may have a write port or communication interface **602** allowing for memory read/write operations. For instance, the communication interface **602** may support packets written to or read from an external memory or an internal static random-access memory (SRAM). The communication interface **602** may employ any suitable protocol such as advanced extensible interface (AXI) protocol. AXI is a high-speed/high-end on-chip bus protocol and has channels associated with read, write, address, and write response, which are respectively separated, individually operated, and have transaction properties such as multiple-outstanding address or write data interleaving. The AXI interface **602** may include features that support unaligned data transfers using byte strobes, burst based transactions with only start address issued, separate address/control and data phases, issuing of multiple outstanding addresses with out of order responses, and easy addition of register stages to provide timing closure. For example, when the MPU executes a table write instruction, the MPU may track which bytes have been written to (a.k.a. dirty bytes) and which remain unchanged. When the table entry is flushed back to the memory, the dirty byte vector may be provided to AXI as a write strobe, allowing multiple writes to safely update a single table data structure as long as they do not write to the same byte. In some cases, dirty bytes in the table need not be contiguous and the MPU may only write back a table if at least one bit in the dirty vector is set. Though packet data is transferred according the AXI protocol in the on-chip communications fabric system according to the examples in the present specification, it may also be applied to a packet data communication on-chip interconnect system operating by other protocols supporting a lock operation, such as advanced high-performance bus (AHB) protocol or advanced peripheral bus (APB) protocol in addition to the AXI protocol.

(63) The MPU **601** may have an instruction fetch unit **605** configured to fetch instructions from a memory external to the MPU based on the table lookup result or at least a portion of the table lookup result. The instruction fetch unit may support branches and/or linear code paths based on table results or a portion of a table result provided by a table engine. In some cases, the table result may comprise table data, key data and/or a start address of a set of instructions/program. The instruction fetch unit **605** may have an instruction cache **604** for storing one or more programs. In some cases, the one or more programs may be loaded into the instruction cache **604** upon receiving the start address of the program provided by the table engine. In some cases, a set of instructions or a program may be stored in a contiguous region of a memory unit, and the contiguous region may be identified by the address. In some cases, the one or more programs may be fetched and loaded from an external memory via the communication interface **602**. This provides flexibility to allow for executing different programs associated with different types of data using the same processing unit. In an example, a management PHV may be injected into the pipeline, for example to perform administrative table direct memory access (DMA) operations or entry aging functions (i.e., adding timestamps), one of the management MPU programs may be loaded to the instruction cache to execute the management function. The instruction cache **604** may be implemented using various types of memories such as one or more SRAMs.

(64) The one or more programs may be any programs such as P4 programs related to reading table

data, building headers, DMA to/from memory, writing to/from memory, and various other actions. The one or more programs may be executed in any match-action processing stage.

(65) The MPU **601** may have a register file unit **606** to stage data between the memory and the functional units of the MPU, or between the memory external to the MPU and the functional units of the MPU. The functional units may include, for example, ALUs, meters, counters, adders, shifters, edge detectors, zero detectors, condition code registers, status registers, and the like. In some cases, the register file unit **606** may comprise a plurality of general-purpose registers (e.g., R0, R1, . . . Rn) which may be initially loaded with metadata values then later used to store temporary variables within execution of a program until completion of the program. For example, the register file unit **606** may be used to store SRAM addresses, ternary content addressable memory (TCAM) search values, ALU operands, comparison sources, or action results. The register file unit of a stage may also provide data/program context to the register file of the subsequent stage, as well as making data/program context available to the next stage's execution data path (i.e., the source registers of the next stage's adder, shifter, and the like). In some MPUs, each register of the register file is 64 bits and may be initially loaded with special metadata values such as hash value from table lookup, packet size, PHV timestamp, programmable table constant and the like.

(66) The register file unit **606** may have a comparator flags unit (e.g., C0, C1, . . . Cn) configured to store comparator flags. The comparator flags may be set by calculation results generated by the ALU which in return may be compared with constant values in an encoded instruction to determine a conditional branch instruction. The MPU may have one-bit comparator flags (e.g., 8 one-bit comparator flags). In practice, an MPU may have any number of comparator flag units, each of which may have any suitable length.

(67) The MPU **601** may have one or more functional units such as the ALU(s) **607**. An ALU may support arithmetic and logical operations on the values stored in the register file unit **606**. The results of the ALU operations (e.g., add, subtract, AND, OR, XOR, NOT, AND NOT, shift, and compare) may then be written back to the register file. The functional units of the MPU may, for example, update or modify fields anywhere in a PHV, write to memory (e.g., table flush), or perform operations that are not related to PHV update. For example, an ALU may be configured to perform calculations on descriptor rings, scatter gather lists (SGLs), and control data structures loaded into the general purpose registers from the host memory.

(68) The MPU **601** may have other functional units such as meters, counters, action insert units, and the like. For example, an ALU may be configured to support P4 compliant meters. A meter is a type of action executable on a table match used to measure data flow rates. A meter may include a number of bands, typically two or three, each of which has a defined maximum data rate and optional burst size. Using a leaky bucket analogy, a meter band is a bucket filled by the packet data rate and drained at a constant allowed data rate. Overflow occurs if the integration of data rate exceeding quota is larger than the burst size. Overflowing one band triggers activity into the next band, which presumably allows a higher data rate. In some cases, a field of the packet may be marked as a result of overflowing the base band. This information might be used later to direct the packet to a different queue, where it may be more subject to delay or dropping in case of congestion. The counter may be implemented by the MPU instructions. The MPU may have one or more types of counters for different purposes. For example, the MPU may have performance counters to count MPU stalls. An action insert unit or set of instructions may be configured to push the register file result back to the PHV for header field modifications.

(69) The MPU may be capable of locking a table. In some cases, a table being processed by an MPU may be locked or marked as "locked" in the table engine. For example, while an MPU has a table loaded into its register file, the table address may be reported back to the table engine, causing future reads to the same table address to stall until the MPU has released the table lock. For instance, the MPU may release the lock when an explicit table flush instruction is executed, the MPU program ends, or the MPU address is changed. In some cases, an MPU may lock more than

one table address, for example, one for the previous table write-back and another address lock for the current MPU program.

(70) A single MPU may be configured to execute instructions of a program until completion of the program. Multiple MPUs may be configured to execute a program. A table result may be distributed to multiple MPUs. The table result may be distributed to multiple MPUs according to an MPU distribution mask configured for the tables. This provides advantages to prevent data stalls or mega packets per second (MPPS) decrease when a program is too long. For example, if a PHV requires four table reads in one stage, then each MPU program may be limited to only eight instructions in order to maintain a 100 MPPS if operating at a frequency of 800 MHz in which scenario multiple MPUs may be desirable.

(71) FIG. 7 illustrates a block diagram of a packet processing pipeline circuit **701** that may be included in the exemplary system of FIG. 4. The packet processing pipeline circuit **701** may be a P4 pipeline circuit in a semiconductor chip. The packet processing pipeline circuit **701** may be programmed to provide various features, including, but not limited to, routing, bridging, tunneling, forwarding, network ACLs, layer 4 firewalls, flow based rate limiting, VLAN tag policies, membership, isolation, multicast and group control, label push/pop operations, layer 4 load balancing, layer 4 flow tables for analytics and flow specific processing, DDOS attack detection, mitigation, telemetry data gathering on any packet field or flow state and various others.

(72) A programmer or compiler may decompose a packet processing program or flow processing data into a set of dependent or independent table lookup and action processing stages (i.e., match-action) that may be mapped onto the table engine and MPU stages. The match-action pipeline may have a plurality of stages. For example, a packet entering the pipeline may be first parsed by a parser stage (e.g., parser **704**) according to the packet header stack specified by a P4 program. This parsed representation of the packet may be referred to as a packet header vector (PHV). The PHV may then be passed through match-action processing stages (e.g., match-action processing stages **705**, **710**, **711**, **712**, **713**, **714**) of the match-action pipeline. Each match-action processing stage may be configured to match one or more PHV fields to tables and to update the PHV, table entries, or other data according to the actions specified by the P4 program. If the required number of stages exceeds the implemented number of stages, a packet may be recirculated for additional processing. The packet payload may travel in a separate queue or buffer until it is reassembled with its PHV in a deparser **715**. The deparser **715** may rewrite the original packet according to the PHV fields which may have been modified in the pipeline. A packet processed by an ingress pipeline may be placed in a packet buffer for scheduling and possible replication. In some cases, once the packet is scheduled and leaves the packet buffer, it may be parsed again to create an egress PHV. The egress PHV may be passed through a P4 egress pipeline in a similar fashion as a packet passing through a P4 ingress pipeline, after which a final deparser operation may be executed before the packet is sent to its destination interface or recirculated for additional processing. The networking device **430** of FIG. 4 may have a P4 pipeline that is implemented via a packet processing pipeline circuit **701**.

(73) A pipeline may have multiple parsers and may have multiple deparsers. The parser may be a P4 compliant programmable parser and the deparser may be a P4 compliant programmable deparser. The parser may be configured to extract packet header fields according to P4 header definitions and place them in a PHV. The parser may select from any fields within the packet and align the information from the selected fields to create the PHV. The deparser may be configured to rewrite the original packet according to an updated PHV. The pipeline MPUs of the match-action processing stages **705**, **710**, **711**, **712**, **713**, **714** may be the same as the MPU **601** of FIG. 6. Match-action processing stages may have any number of MPUs. The match-action processing stages of a match-action pipeline may all be identical.

(74) A table engine **706** may be configured to support per-stage table match. For example, the table engine **706** may be configured to hash, lookup, and/or compare keys to table entries. The table engine **706** may be configured to control the address and size of the table, use PHV fields to

generate a lookup key, and find Session Ids or MPU instruction pointers that define the P4 program associated with a table entry. A table result produced by the table engine may be distributed to the multiple MPUs.

(75) The table engine **706** may be configured to control a table selection. In some cases, upon entering a stage, a PHV is examined to select which table(s) to enable for the arriving PHV. Table selection criteria may be determined based on the information contained in the PHV. In some cases, a match table may be selected based on packet type information related to a packet type associated with the PHV. For instance, the table selection criteria may be based on a debug flag, packet type or protocols (e.g., Internet Protocol version 4 (IPv4), Internet Protocol version 6 (IPv6), MPLSA, or the next table ID as determined by the preceding stage. In some cases, the incoming PHV may be analyzed by the table selection logic, which then generates a table selection key and compares the result using a TCAM to select the active tables. A table selection key may be used to drive table hash generation, table data comparison, and associated data into the MPUs.

(76) The table engine **706** may have a ternary content-addressable memory (TCAM) control unit **708**. The TCAM control unit may be configured to allocate memory to store multiple TCAM search tables. In an example, a PHV table selection key may be directed to a TCAM search stage before a SRAM lookup. The TCAM control unit may be configured to allocate TCAMs to individual pipeline stages to prevent TCAM resource conflicts, or to allocate TCAM into multiple search tables within a stage. The TCAM search index results may be forwarded to the table engine for SRAM lookups.

(77) The table engine **706** may be implemented by hardware or circuitry. The table engine may be hardware defined. In some cases, the results of table lookups or table results are provided to the MPU in its register file.

(78) A match-action pipeline may have multiple match-action processing stages such as the six units illustrated in the example of FIG. 7. In practice, a match-action pipeline may have any number of match-action processing stages. The match-action processing stages may share a pipeline memory circuit **702** that may be static random-access memory (SRAM), TCAM, some other type of memory, or a combination of different types of memory. The packet processing pipeline circuit stores data in the pipeline memory circuit. For example, the packet processing pipeline circuit may store a table in the pipeline memory circuit that configures the packet processing pipeline circuit to process specific network flows. For example, a flow table or multiple flow tables may be stored in the pipeline memory circuit **702** and may store instructions and data that the packet processing pipeline circuit uses to process a packet. The pipeline memory circuit is more than half full when it is storing data used by the packet processing pipeline circuit and less than half the capacity of the pipeline memory circuit is free.

(79) The second match-action pipeline circuit **425** includes a match-action pipeline **434**. That match-action pipeline **434** may include match-action processing stages such as match-action processing stages **705, 710, 711, 712, 713, 714**.

(80) FIG. 8 illustrates packet headers and payloads of packets for a network flow **800** including layer 7 fields according to some aspects. A group of network packets passing from one specific endpoint to another specific endpoint is a network flow. A network flow **800** may have numerous network packets such as a first packet **850**, a second packet **851**, a third packet **852**, a fourth packet **853**, and a final packet **854** with many more packets between the fourth packet **853** and the final packet **854**. The term “the packet” or “a packet” may refer to any of the network packets in a network flow.

(81) Packets may be constructed and interpreted in accordance with the internet protocol suite. The Internet protocol suite is the conceptual model and set of communications protocols used in the Internet and similar computer networks. A packet may be transmitted and received as a raw bit stream over a physical medium at the physical layer, sometimes called layer 1. The packets may be received by a RX MAC **211** as a raw bit stream or transmitted by TX MAC **210** as a raw bit stream.

(82) The link layer is often called layer 2. The protocols of the link layer operate within the scope of the local network connection to which a host is attached and includes all hosts accessible without traversing a router. The link layer is used to move packets between the interfaces of two different hosts on the same link. The packet (an Ethernet packet is shown) has a layer 2 header **801**, a layer 2 payload **802**, and a layer 2 frame check sequence (FCS) **803**. The layer 2 header may contain a source MAC address **804**, a destination MAC address **805**, an optional 802.1Q header **806**, optional VLAN tag information **807**, and other layer 2 header data **808**. The input ports **211** and output ports **210** of a networking device **201** may have MAC addresses. A networking device **201** may have a MAC address that is applied to all or some of the ports. Alternatively, a networking device may have one or more ports that each have their own MAC address. In general, each port may send and receive packets. As such, a port of a networking device may be configured with a RX MAC **211** and a TX MAC **210**. Ethernet, also known as Institute of Electrical and Electronics Engineers (IEEE) 802.3, is a layer 2 protocol. IEEE 802.11 (WIFI) is another widely used layer 2 protocol. The layer 2 payload **802** may include a layer 3 packet. The layer 2 FCS **803** may include a CRC (cyclic redundancy check) calculated from the layer 2 header and layer 2 payload. The layer 2 FCS may be used to verify that the packet has been received without errors.

(83) IEEE 802.1Q is the networking standard that supports VLANs on IEEE 802.3 networks. The optional 802.1Q header **806** and VLAN tag information **807** are specified by the IEEE 802.1Q standard. The 802.1Q header is the two-octet value 0x8100 that indicates that VLAN tag information **807** is present. The VLAN tag information includes a 12-bit VLAN identifier. As such, a LAN may be configured to have 4094 VLANs (0x000 and 0xFFF are reserved values).

(84) The internet layer, often called layer 3, is the network layer where layer 3 packets may be routed from a first node to a second node across multiple intermediate nodes. The nodes may be networking devices such as networking device **201**. Internet protocol (IP) is a commonly used layer 3 protocol that is specified in requests for comment (RFCs) published by the Internet Engineering Task Force (IETF). More specifically, the format and fields of IP packets are specified by IETF RFC 791. The layer 3 packet (an IP packet is shown) may have a layer 3 header **810** and a layer 3 payload **811**. The layer 3 header of an IP packet is an IP header and the layer 3 payload of an IP packet is an IP payload. The layer 3 header **810** may have a source IP address **812**, a destination IP address **813**, a protocol indicator **814**, and other layer 3 header data **815**. In general, a packet is directed from a source machine at the source IP address **812** and is directed to the destination machine at the destination IP address **813**. As an example, a first node may send an IP packet to a second node via an intermediate node. The IP packet therefore has a source IP address indicating the first node and a destination IP address indicating the second node. The first node may make a routing decision to send the IP packet to the intermediate node. The first node may therefore send the IP packet to the intermediate node in a first layer 2 packet. The first layer 2 packet has a source MAC address **804** indicating the first node, a destination MAC address **805** indicating the intermediate node, and has the IP packet as a payload. The intermediate node receives the first layer 2 packet. Based on the destination IP address, the intermediate node may make a routing decision to send the IP packet to the second node. The intermediate node may then send the IP packet to the second node in a second layer 2 packet having a source MAC address **804** indicating the intermediate node, a destination MAC address **805** indicating the second node, and the IP packet as a payload. The layer 3 payload **811** may include headers and payloads for higher layers in accordance with higher layer protocols such as transport layer protocols.

(85) The transport layer, often called layer 4, may establish basic data channels that applications use for task-specific data exchange and may establish host-to-host connectivity. A layer 4 protocol may be indicated in the layer 3 header **810** using protocol indicator **814**. Transmission control protocol (TCP, specified by IETF RFC 793), user datagram protocol (UDP, specified by IETF RFC 768), and internet control message protocol (ICMP, specified by IETF RFC 792) are common layer 4 protocols. TCP is often referred to as TCP/IP. TCP is connection oriented and may provide

reliable, ordered, and error-checked delivery of a stream of bytes between applications running on hosts communicating via an IP network. When carrying TCP data, a layer 3 payload **811** includes a TCP header and a TCP payload. UDP may provide for computer applications to send messages, in this case referred to as datagrams, to other hosts on an IP network using a connectionless model. When carrying UDP data, a layer 3 payload **811** includes a UDP header and a UDP payload. ICMP is used by network devices, including routers, to send error messages and operational information indicating success or failure when communicating with another IP address. ICMP uses a connectionless model.

(86) A layer 4 packet (a TCP packet is shown) may have a layer 4 header **820** (a TCP header is shown) and a layer 4 payload **821** (a TCP payload is shown). The layer 4 header **820** may include a source port **822**, destination port **823**, layer 4 flags **824**, and other layer 4 header data **825**. The source port and the destination port may be integer values used by host computers to deliver packets to application programs configured to listen to and send on those ports. The layer 4 flags **824** may indicate a status of or action for a network flow. A layer 4 payload **821** may contain a layer 7 packet.

(87) The application layer, often called layer 7, includes the protocols used by most applications for providing user services or exchanging application data over the network connections established by the lower level protocols. Examples of application layer protocols include NVMe/TCP, RDMA over Converged Ethernet version 2, (RoCE v2), Hypertext Transfer Protocol (HTTP), File Transfer Protocol (FTP), Simple Mail Transfer Protocol (SMTP), and Dynamic Host Configuration (DHCP). Data coded according to application layer protocols may be encapsulated into transport layer protocol data units (such as TCP or UDP messages), which in turn use lower layer protocols to effect actual data transfer.

(88) A layer 4 payload **821** may include a layer 7 packet **830**. The illustrated layer 7 packet is a NVMe/TCP PDU **830**. NVM Express, Inc. is a trade organization that provides standards for NVMe communications. Version 1.0 of the NVM Express TCP Transport Specification was published on May 18, 2021 and discloses the contents and structure of NVMe PDUs. An NVMe/TCP PDU may include a header **831**, a header digest field **832**, a pad **834**, a PDU data field **835**, and a data digest field **836**. The PDU data field **835** may hold numerous encrypted blocks. When the PDU data field **835** holds numerous encrypted blocks, the data digest field **836** may hold the digest value for the entire PDU data field **835**. Such a digest value may be determined from the individual digest values of the numerous encrypted blocks in the PDU data field.

(89) FIG. 9 illustrates a high-level flow diagram of a process **900** that may be implemented by packet processing pipeline obtaining a policing decision from a programmable policer circuit, according to some aspects. After the start, a network packet in a network flow is received at block **901**. At block **902**, a PHV is produced by parsing the network packet. At block **903**, the PHV is used to identify a flow table entry. The network packet is in a network flow and the flow table entry may be used for processing all the network packets in the network flow. For example, the workload (e.g., a VM or other network endpoint) may send a network packet in a network flow that is a storage flow. All the network packets in a network flow may have the same 5-tuple. The network flow may be identified as a storage flow from the 5-tuple, the presence of a NVMe/TCP PDU in the network packets of the network flow, etc. Those practiced in computer networking realize that TCP ports **8009** and **4420** are default ports for NVMe/TCP and that other storage protocols may use known default ports, or protocol identifiers. In another example, the workload may send a network packet in a network flow that is a non-storage flow. The network flow may be identified as a non-storage flow from the 5-tuple, the presence of HTTP or HTTPS headers or data in the network packets of the network flow, etc. Those practiced in computer networking realize that port **80** is the default port for HTTP, port **443** is the default port for HTTPS, and that other traffic may also have default ports. In general, the type of a network flow may be determined from the 5-tuple and other PHV contents.

(90) At block **904**, the programmable policer identifier and the policing policy identifier may be read from the flow table entry. The programmable policer identifier identifies one of the programmable policer circuits. At block **905**, policer data may be sent to the programmable policer circuit. The policer data may include the policing policy identifier and a resources required value (e.g., the packet size in bytes). At block **906**, a policing decision may be received from the programmable policer circuit. At decision block **907**, the process may determine if the policing decision is to drop the packet. If the policing decision is to drop the packet, the process may move to block **908**. Otherwise, the process may move to block **909**. At block **908**, the packet is dropped and the process is done. At block **909**, the packet processing pipeline circuit continues processing the network packet and may produce a processed network packet that is transmitted from the egress port of the networking device.

(91) FIG. **10** is a high-level conceptual diagram illustrating a metering circuit refilling the tokens in a token bucket **1001** that may be used as an aggregated token bucket according to some aspects. A semiconductor chip such as the semiconductor chip **401** illustrated in FIG. **4** may include the metering circuit **423**. The token bucket in the illustrated example is a leaky token bucket. The token bucket **1001** may include a tokens value **141**, a token add rate value **1003**, and a maximum tokens value **1004**. The tokens value **141** indicates how many tokens are in the bucket. The token add rate value **1003** indicates the number of tokens to be added to the tokens value **141** every metering interval. The max tokens value **1004** indicates the maximum number of tokens that may be in the token bucket **1001**. The tokens value may be a signed integer. Max tokens **1004** and the token add rate **1003** may be set based on the QoS that is to be provided to the workload or tenant running the workload. An interval timer **1005** may trigger the metering circuit **423** every metering interval (e.g., every 0.1 seconds). Every time it is triggered, the metering circuit **423** may add the number of tokens indicated by the token add rate **1003** to the tokens value **141** up to the maximum indicated by the maximum token value **1004**. The tokens may be removed from the token bucket by a programmable policer circuit **123**. A number of tokens are removed from the token bucket by subtracting that number from the tokens value. For example, 10 tokens are removed by subtracting 10 from the tokens value. For example, in FIG. **1** the programmable policer circuit **123** may remove tokens by subtracting the packet size **121** from the tokens value **141**.

(92) FIG. **11** is a high-level conceptual diagram illustrating policing policy identifiers **1102** that are associated with policing policies **1103**, according to some aspects. The policies identified with the identifiers are presented in the example of FIG. **11** as a policer policy table **1101**. In many implementations, there is no explicit table. For example, the programmable policer circuit may be a circuit that includes hardware blocks that perform different policing policies where one of the hardware blocks is selected by the value of the policing policy identifier. The programmable policer circuit may implement a policing policy that produces a “drop” policing decision when the policing policy identifier equals 0 and allowing the packet would cause the tokens value to be less than a threshold value. For example, a value equaling the tokens value minus the resources required (e.g., packet size) may be stored in a first register of the programmable policer circuit and the threshold value may be stored in a second register of the policer circuit. If the value in the first register is less than the value in the second register, then allowing the packet may cause the tokens value to be less than a threshold value. If allowing the packet would not cause the tokens value to be less than a threshold value, then the programmable policer circuit may produce an “allow” policing decision and subtract the resources required value from the tokens value when the policing policy identifier equals 0. The threshold value is often zero, in which case a negative tokens value is less than the threshold value. The programmable policer circuit may implement a policing policy that always produces an “allow” policing decision and always subtracts the resources required value from the tokens value when the policing policy identifier equals 1. This policy would allow the packet even though allowing the packet would cause the tokens value to be less than a threshold value (e.g., less than 0). When the tokens value is less than the threshold value, the network packets

for network flows that have a policing policy equaling 0 will always be dropped. No such network packet will be allowed until the tokens value equals or exceeds the threshold value. The tokens value may eventually become greater than the threshold value as the metering circuit **423** adds tokens to the token bucket. As such, the network flows having a policing policy equaling 1 are strictly prioritized over network flows having a policing policy equaling 0. The programmable policer circuit may implement a policing policy that produces an “allow” policing decision without reading or changing the tokens value when the policing policy identifier equals 2. Such a policy may be used within the data center for the network flows of infrastructure operations such as syncing flow tables between networking devices, configuration changes, etc.

(93) FIG. **12** is a high-level flow diagram illustrating an example of a process for producing policing decisions and that may be implemented by a programmable policer circuit, according to some aspects. After the start, policing data is received at block **1201**. The policing data may include a policing policy ID and a resource required value such as the packet size in bytes. At decision block **1202**, the process determines if the policing policy identifier equals 0. If the policing policy identifier equals 0 at decision block **1202**, the process moves to block **1203** and otherwise moves to decision block **1209**. At block **1203**, the process may calculate a new value that equals the resources required value subtracted from the tokens value. At decision block **1204**, the process determines if the new value is less than the threshold value (e.g., new value is negative when the threshold value is 0). If the new value is less than the threshold value at decision block **1204**, the process moves to block **1205** and otherwise moves to decision block **1206**. At block **1205**, the policing decision is “drop”. The policing decision is transmitted to the packet processing pipeline circuit at block **1208**. At block **1206**, the policing decision is “allow”. At block **1207**, the tokens value is set to the new value calculated at block **1203**. The policing decision is transmitted to the packet processing pipeline circuit at block **1208**.

(94) At decision block **1209**, the process determines if the policing policy identifier equals 1. If the policing policy identifier equals 1 at decision block **1209**, the process moves to block **1210** and otherwise moves to decision block **1213**. At block **1210**, the process may decrease the tokens value by subtracting the resources required value from the tokens value. At block **1211**, the policing decision is “allow”. The policing decision is transmitted to the packet processing pipeline circuit at block **1212**. At decision block **1213**, the process determines if the policing policy identifier equals 2. If the policing policy identifier equals 2 at decision block **1213**, the process moves to block **1214** and otherwise moves to block **1216**. At block **1214**, the policing decision is “allow”. The policing decision is transmitted to the packet processing pipeline circuit at block **1215**. If the process reaches block **1216**, an error has occurred such as when there is no policing policy associated with the policing policy identifier. At block **1216**, the process may perform a default operation such as raising an error flag or signaling an error.

(95) FIG. **13** is a high-level conceptual diagram illustrating an example of allowing a packet received from a network endpoint **1301**, according to some aspects. The network endpoint **1301** may be a VM attempting to send a storage packet **1302**. A storage packet is a network packet in a storage flow. The parser **1304** parses the storage packet **1302** to produce a PHV **1305**. The first processing stage **1306** reads the flow table entry **1321** for the storage flow from the flow table **1320**. In this example, the programmable policer identifier is 55 and the policing policy identifier is 1. The first processing stage adds the programmable policer identifier and the policing policy identifier to the PHV **1307**. The second processing stage **1308** receives the PHV **1307** and sends policing data **1309** to the 55th programmable policer circuit **1311**, which is the programmable policer circuit identified by the programmable policer identifier. The policing data **1309** includes the packet size (the resources required value) and the policing policy identifier. The 55th programmable policer circuit **1311** is associated with the 28th aggregated token bucket **1313**. For example, the programmable policer circuits may contain an aggregated token bucket identifier such as the memory address of the aggregated token bucket. The aggregated token bucket identifier

1312 of the 55th programmable policer circuit **1311** may identify the 28th aggregated token bucket **1313**. In this example, the tokens value **1314** in the 28th aggregated token bucket **1313** is less than the packet size. When the policing policy identifier equals 1, the 55th programmable policer circuit **1311** produces an “allow” policing decision **1310** and subtracts the resources required value from the tokens value. The second processing stage receives the policing decision and passes the PHV **1315** to the other stages **1316** of the packet processing pipeline for further processing of the storage packet **1302** when the policing decision **1310** is “allow”.

(96) FIG. **14** is a high-level conceptual diagram illustrating an example of dropping a packet received from a network endpoint **1301**, according to some aspects. The network endpoint **1301** attempts to send a non-storage packet **1303**. A non-storage packet is a network packet in a non-storage flow. The parser **1304** parses the non-storage packet **1303** to produce a PHV **1401**. The first processing stage **1306** reads the flow table entry **1322** for the non-storage flow from the flow table **1320**. In this example, the programmable policer identifier is 55 and the policing policy identifier is 0. The first processing stage adds the programmable policer identifier and the policing policy identifier to the PHV **1402**. The second processing stage **1308** receives the PHV **1402** and sends policing data **1403** to the 55th programmable policer circuit **1311**, which is the programmable policer circuit identified by the programmable policer identifier. The policing data **1309** includes the packet size (the resources required value) and the policing policy identifier. The 55th programmable policer circuit **1311** is associated with the 28th aggregated token bucket **1313**. For example, the programmable policer circuits may contain (e.g., store in a register) an aggregated token bucket identifier. The aggregated token bucket identifier **1312** of the 55th programmable policer circuit **1311** may identify the 28th aggregated token bucket **1313**. In this example, the tokens value **1314** in the 28th aggregated token bucket **1313** is less than the packet size. The 55th programmable policer circuit **1311** produces a “drop” policing decision **1404** when the policing policy identifier equals 0 and the resources required value is greater than the tokens value (assumes threshold equals 0). The second processing stage receives the policing decision and drops the non-storage packet **1303**.

(97) FIG. **15** is a high-level conceptual diagram illustrating a control plane **1502** using networking policies **1501** to produce flow table entries, according to some aspects. The networking policies **1501** and a flow table **110** that includes flow table entries may be stored in a memory of a networking device. A network packet is received and parsed to produce a PHV. The match unit **1505** of a processing stage **1504** attempts to read a flow table entry for the network packet from the flow table **110**. The match unit may produce a flow miss **1506** when the flow entry does not exist. The action unit **1507** may add a flow miss indicator to the PHV and send the PHV to the control plane **1502**. A flow table entry compiler **1503** may produce flow table entries for the network packet by applying the networking policies to the PHV. The flow table entries thus produced may be specifically for processing the network packets of the network flow that includes the network packet that caused the flow miss. A flow table entry **1508** for the network flow may then be stored in the flow table. Other flow table entries may be stored in the flow tables that are used by other processing stages. The data plane, which includes the packet processing pipeline circuit, may be configured to process a network flow when the control plane stores the flow table entries for the network flow in the flow tables.

(98) FIG. **16** is a high-level flow diagram illustrating a method **1600** for using programmable policer circuits to implement network flow policing, according to some aspects. At block **1601**, the method may use a plurality of networking policies to produce a plurality of flow table entries that includes the flow table entry. At block **1602**, the method may associate a programmable policer circuit and a first policing policy with a first network flow. For example, a flow table entry for the first network flow may be created that includes the policer circuit ID of the programmable policer circuit and the policing policy ID of the first policing policy. At block **1603**, the method may associate the programmable policer circuit and a second policing policy with a second network

flow. For example, a flow table entry for the second network flow may be created that includes the policer circuit ID of the programmable policer circuit and the policing policy ID of the second policing policy. At block **1604**, the method may store a tokens value of an aggregated token bucket. At block **1605**, a packet processing pipeline circuit, may use a first network packet that is in the first network flow to identify the programmable policer circuit and the first policing policy. At block **1606**, the programmable policer circuit may use the first policing policy and the tokens value to produce a first policing decision for the first network packet. At block **1607**, the packet processing pipeline circuit may use a second network packet that is in the second network flow to identify the programmable policer circuit and the second policing policy. At block **1608**, the programmable policer circuit may use the second policing policy to produce a second policing decision for the first network packet. At block **1609**, the method may process the first network packet in accordance with the first policing decision. At block **1610**, the method may process the second network packet in accordance with the second policing decision. A packet may be processed in accordance with a policing decision by implementing the policing decision while processing the packet. For example, a packet may be processed in accordance with a “drop” policing decision by dropping the packet. A packet may be processed in accordance with an “allow” policing decision by not dropping the packet.

(99) FIG. **17** is another high-level flow diagram illustrating a method **1700** for using programmable policer circuits to implement network flow policing, according to some aspects. At block **1701**, the method may use a network packet to identify a programmable policer circuit and a policing policy for a network flow that includes the network packet. At block **1702**, the method may obtain a policing decision for the network packet from the programmable policer circuit, the policing decision in accordance with the policing policy. At block **1703**, a packet processing pipeline circuit may process the network packet in accordance with the policing decision.

(100) Aspects described above may be ultimately implemented in a networking device that includes physical circuits that implement digital data processing, storage, and communications. The networking device may include processing circuits, ROM, RAM, TCAM, and at least one interface (interface(s)). The CPU cores described above are implemented in processing circuits and memory that is integrated into the same integrated circuit (IC) device as ASIC circuits and memory that are used to implement the programmable packet processing pipeline. For example, the CPU and other semiconductor chip circuits are fabricated on the same semiconductor substrate to form a System-on-Chip (SoC). The networking device may be implemented as a single IC device (e.g., fabricated on a single substrate) or the networking device may be implemented as a system that includes multiple IC devices connected by, for example, a printed circuit board (PCB). The interfaces may include network interfaces (e.g., Ethernet interfaces and/or InfiniBand interfaces) and/or PCIe interfaces. The interfaces may also include other management and control interfaces such as I2C, general purpose IOs, USB, UART, SPI, and eMMC.

(101) FIG. **18** is a high-level block diagram illustrating an example of a programmable policer circuit **123**, according to some aspects. The programmable policer circuit **123** may include policer circuit registers **1801**, a policer logic selection circuit **1802**, policer logic **1803**, and a token bucket interface **1807**. The policer circuit registers **1801** may include a policing policy identifier register, a requested amount register, and a token bucket indicator register. The policing policy identifier **122** in the policing data may be written into the policy identifier register. The packet size value **121** in the policing data **120** may be written into the requested amount register. A token bucket may be associated with the programmable policer circuit by writing a token bucket indicator into the token bucket ID register. The token bucket interface **1807** may use the token bucket indicator to access the token bucket **1001** to, for example, read and write the tokens value **141**. The token bucket indicator may be an index value (e.g., “33” to indicate the 33rd token bucket), an address, or any other value that identifies or indicates a token bucket. The control plane may write a value into the token bucket indicator register during initialization of the data plane. The flow table entry for a

network flow may include a token bucket indicator that may be included in the policing data **120** and that may be written into the token bucket indicator register when the other values in the policing data **120** are written into the policer circuit registers. In such a scenario, a token bucket is associated with a network flow when the control plane configures the data plane to process network packets for the network flow. The token bucket **1001** is then associated with the programmable policer circuit **123** when the policing data **120** is written into the policer circuit registers **1801**. A value may be hardwired into the token bucket indicator register such that a specific token bucket is always associated with a specific programmable policer circuit.

(102) The policer logic circuits **1803** may produce policing decisions **124**. The policer logic circuits **1803** may include a first policing policy logic circuit **1804**, a second policing policy logic circuit **1805**, and a third policing policy logic circuit **1806**. The policer logic selection circuit **1802** may use the value in the policer policy identifier register to select one of the policer logic circuits **1803** to thereby select a policing policy. For example, the first policing policy logic circuit **1804** may produce a “drop” policing decision in response to determining that the tokens value **141** minus the requested amount is less than a threshold value. Otherwise, the first policing policy logic circuit **1804** may produce an “allow” policing decision and subtract the requested amount from the tokens value in response to determining that the tokens value **141** minus the requested amount is not less than a threshold value. The second policing policy logic circuit **1805** may always produce an “allow” policing decision and always subtract the requested amount from the tokens value. The third policing policy logic circuit **1806** may always produce an “allow” policing decision and not subtract the requested amount from the tokens value.

(103) Although the operations of the method(s) herein are shown and described in a particular order, the order of the operations of each method may be altered so that certain operations may be performed in an inverse order or so that certain operations may be performed, at least in part, concurrently with other operations. Instructions or sub-operations of distinct operations may be implemented in an intermittent and/or alternating manner.

(104) It may also be noted that at least some of the operations for the methods described herein may be implemented using software instructions stored on a computer usable storage medium for execution by a computer. For example, a computer program product may include a computer usable storage medium to store a computer readable program.

(105) The computer-usable or computer-readable storage medium may be an electronic, magnetic, optical, electromagnetic, infrared, or semiconductor system (or apparatus or device). Examples of non-transitory computer-usable and computer-readable storage media include a semiconductor or solid-state memory, magnetic tape, a removable computer diskette, a random-access memory (RAM), a read-only memory (ROM), a rigid magnetic disk, and an optical disk. Current examples of optical disks include a compact disk with read only memory (CD-ROM), a compact disk with read/write (CD-R/W), and a digital video disk (DVD).

(106) Although specific examples have been described and illustrated, the scope of the claimed systems, methods, devices, etc. is not to be limited to the specific forms or arrangements of parts so described and illustrated. The scope is to be defined by the claims appended hereto and their equivalents.

Claims

1. A system comprising: a programmable policer circuit configured to use a token bucket and a policing policy to produce a policing decision for a network packet in a network flow associated with the policing policy; a packet processing pipeline circuit configured to implement a data plane; and a processor configured to implement a control plane; wherein the programmable policer circuit, the packet processing pipeline circuit, and the processor are further configured to implement network flow policing that involves: the control plane configuring the data plane to process the

network packet; and the control plane associating the network flow with the policing policy such that the data plane is configured to process the network packet in accordance with the policing decision.

2. The system of claim 1, wherein: the packet processing pipeline circuit is configured to use the network packet to identify a flow table entry that includes a policing policy identifier identifying the policing policy and a policer identifier identifying the programmable policer circuit; and the packet processing pipeline circuit uses the policer identifier to provide policing data to the programmable policer circuit, the policing data including the policing policy identifier.

3. The system of claim 2 wherein the control plane uses a plurality of networking policies to produce a plurality of flow table entries that includes the flow table entry.

4. The system of claim 2, wherein: the packet processing pipeline circuit includes a parser stage and a plurality of match-action processing stages; the parser stage produces a packet header vector (PHV) from the network packet; and one of the match-action processing stages uses the PHV to identify the flow table entry.

5. The system of claim 1, wherein: the token bucket is an aggregated token bucket; the programmable policer circuit is configured to use the aggregated token bucket and a second policing policy to produce a second policing decision for a second network packet in a second network flow associated with the second policing policy; and the packet processing pipeline circuit is configured to process the second network packet in accordance with the second policing decision.

6. The system of claim 5 wherein the network flow is directed to or from a network endpoint and the second network flow is directed to or from the network endpoint.

7. The system of claim 6 wherein the network endpoint is a virtual machine.

8. The system of claim 5 wherein: the policing decision is to drop the network packet in response to determining that allowing the network packet would cause a tokens value of the aggregated token bucket to be less than a threshold value; and the second policing decision is to allow the second network packet even though the tokens value is below the threshold value.

9. The system of claim 8 wherein: the network flow is a non-storage flow; and the second network flow is a storage flow.

10. The system of claim 5 wherein: the policing decision is to drop the network packet in response to determining that allowing the network packet would cause a tokens value of the aggregated token bucket to be less than a threshold value; and the second policing policy is to always allow and to always decrease the tokens value based on a size of an allowed network packet.

11. A method comprising: using a network packet to identify a programmable policer circuit and a policing policy for a network flow that includes the network packet; obtaining a policing decision for the network packet from the programmable policer circuit, the policing decision in accordance with the policing policy; and processing, by a packet processing pipeline circuit, the network packet in accordance with the policing decision.

12. The method of claim 11, wherein: a flow table includes a flow table entry that is associated with first network flow, the packet processing pipeline circuit uses the network packet to identify the flow table entry; the flow table entry includes a policing policy identifier that identifies the policing policy and includes a policer identifier that identifies the programmable policer circuit; and the packet processing pipeline circuit uses the policer identifier and the policing policy identifier to obtain the policing decision from the programmable policer circuit.

13. The method of claim 12 further including using a plurality of networking policies to produce a plurality of flow table entries that includes the flow table entry.

14. The method of claim 12, wherein: the packet processing pipeline circuit includes a parser stage and a plurality of match-action processing stages; the parser stage produces a packet header vector (PHV) from the network packet; and one of the match-action processing stages uses the PHV to identify the flow table entry.

15. The method of claim 11, further including: using a tokens value in an aggregated token bucket for policing the network flow and a second network flow that is associated with a second policing policy; using, by the programmable policer circuit, the aggregated token bucket and the second policing policy to produce a second policing decision for a second network packet that is in the second network flow; and processing, by the packet processing pipeline circuit, the second network packet in accordance with the second policing decision.
16. The method of claim 15, wherein: the policing decision is to drop the network packet in response to determining that allowing the network packet would cause the tokens value to be less than a threshold value; and the second policing decision is to allow the second network packet even though the tokens value is below the threshold value.
17. The method of claim 15 wherein: the network flow is a non-storage flow for a virtual machine; and the second network flow is a storage flow for the virtual machine.
18. The method of claim 15 wherein: the policing decision is to drop the network packet in response to determining that allowing the network packet would cause the tokens value to be less than a threshold value; and the second policing policy is to always allow and to always decrease the tokens value based on a size of an allowed network packet.
19. A system comprising: a programmable policing means for producing a policing decision for a network packet in accordance with a tokens value and a policing policy for a network flow; and a packet processing means for processing the network packet in accordance with the policing decision, wherein: the programmable policing means produces a second policing decision for a second network packet in accordance with the tokens value and a second policing policy for a second network flow; and the packet processing means processes the second network packet in accordance with the second policing decision.
20. The system of claim 19, further including: a lookup means for using the network packet to identify the programmable policing means and the policing policy, wherein the lookup means uses the second network packet to identify the programmable policing means and the second policing policy.
-