

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388795
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Loman; Mark Willem et al.

Encrypted cache protection

Abstract

Secrets such as secure session cookies for a web browser can be protected on a compute instance with multiple layers of encryption, such as by encrypting key material that in turn controls cryptographic access to the secret. A compute instance can be instrumented to detect when a process attempts to decrypt this key material so that the process requesting decryption can be compared to authorized or legitimate users of the secret.

Inventors: Loman; Mark Willem (Delden, NL), Engels; Lute Edwin (Zuidwolde, NL), Tijink; Ronny Henk Gert (Hengelo, NL), van Hillo; Victor Marinus Johann Simon (Delden, NL), Vermaning; Alexander (Enschede, NL), Harmsen; Jeroen (Hengelo, NL)

Applicant: Sophos Limited (Abingdon, GB)

Family ID: 83448486

Assignee: Sophos Limited (Abingdon, GB)

Appl. No.: 18/600767

Filed: March 10, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240223537 A1	Jul. 04, 2024

Related U.S. Application Data

continuation parent-doc US 17467733 20210907 US 11929992 child-doc US 18600767
us-provisional-application US 63168654 20210331

Publication Classification

Int. Cl.: H04L29/00 (20060101); H04L9/40 (20220101)

U.S. Cl.:

CPC H04L63/0414 (20130101); H04L63/0435 (20130101); H04L63/102 (20130101); H04L63/105 (20130101); H04L63/20 (20130101);

Field of Classification Search

CPC: H04L (63/0414); H04L (63/0435); H04L (63/102); H04L (63/105); H04L (63/20); H04L (2463/061); H04L (2463/062); H04L (67/146); H04L (9/0822); H04L (9/088); H04L (9/0894); G06F (21/602); G06F (21/6209)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8181244	12/2011	Boney	N/A	N/A
8201243	12/2011	Boney	N/A	N/A
8261087	12/2011	Fort et al.	N/A	N/A
8418250	12/2012	Morris et al.	N/A	N/A
8719932	12/2013	Boney	N/A	N/A
8726389	12/2013	Morris et al.	N/A	N/A
8763123	12/2013	Morris et al.	N/A	N/A
8856505	12/2013	Schneider	N/A	N/A
9413721	12/2015	Morris et al.	N/A	N/A
9578045	12/2016	Jaroch et al.	N/A	N/A
10257224	12/2018	Jaroch et al.	N/A	N/A
10262141	12/2018	Kocher et al.	N/A	N/A
10284591	12/2018	Giuliani et al.	N/A	N/A
10785020	12/2019	Balasubramanian et al.	N/A	N/A
10862883	12/2019	Bhattacharyya	N/A	H04L 63/0861
10999289	12/2020	Shah et al.	N/A	N/A
11429658	12/2021	Richter	N/A	G06F 16/483
2003/0041263	12/2002	Devine et al.	N/A	N/A
2003/0182559	12/2002	Curry et al.	N/A	N/A
2005/0195975	12/2004	Kawakita	N/A	N/A
2005/0249225	12/2004	Singhal	N/A	N/A
2007/0124482	12/2006	Lee et al.	N/A	N/A
2007/0245137	12/2006	Bhagat et al.	N/A	N/A
2012/0151118	12/2011	Flynn	711/170	G06F 11/1666
2013/0097369	12/2012	Talagala	711/103	G06F 12/0804

2015/0294123	12/2014	Oxford	N/A	N/A
2017/0041296	12/2016	Ford	N/A	G06F 21/64
2017/0244566	12/2016	Tschache	N/A	N/A
2018/0046635	12/2017	Nichols	N/A	N/A
2018/0102902	12/2017	Yang et al.	N/A	N/A
2019/0005576	12/2018	Mick	N/A	H04L 41/5054
2019/0272388	12/2018	Tsou et al.	N/A	N/A
2020/0249912	12/2019	Ishikawa et al.	N/A	N/A
2021/0320791	12/2020	Lanc et al.	N/A	N/A
2022/0321540	12/2021	Loman et al.	N/A	N/A
2023/0327855	12/2022	Li	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
WO-2022/208045	12/2021	WO	N/A

OTHER PUBLICATIONS

USPTO, “U.S. Appl. No. 17/467,733 Notice of Allowance mailed Nov. 8, 2023”, 16 pages. cited by applicant

WIPO, “PCT Application No. PCT/GB22/050393 International Preliminary Report on Patentability mailed Oct. 12, 2023”, 10 pages. cited by applicant

ISA/EP, “PCT Application No. PCT/GB22/050393 International Search Report and Written Opinion mailed Apr. 22, 2022”, 14 pages. cited by applicant

Primary Examiner: Rahman; Mahfuzur

Attorney, Agent or Firm: Strategic Patents, P.C.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 17/467,733 filed Sep. 7, 2021, which claims priority to U.S. Provisional Patent Application No. 63/168,654 filed on Mar. 31, 2021, where each of the foregoing applications are hereby incorporated by reference in their entirety.

BACKGROUND

(1) Computing devices often contain valuable information such as the credentials stored in a browser's cache, or corresponding cookies and other tokens used to automatically log in to websites. Malicious actors frequently target these potentially valuable assets in an attempt to lift login capabilities or control the cryptographic tools protecting them. There remains a need for improved techniques to detect and prevent attempted access to local secrets such as encrypted password and cookie caches.

SUMMARY

(2) Secrets such as secure session cookies can be protected on a compute instance with multiple layers of encryption, such as by encrypting key material that in turn controls cryptographic access to the secret. A compute instance can be instrumented to detect when a process attempts to decrypt

this key material so that the process requesting decryption can be compared to authorized or legitimate users of the secret.

Description

BRIEF DESCRIPTION OF THE FIGURES

(1) The foregoing and other objects, features and advantages of the devices, systems, and methods described herein will be apparent from the following description of particular embodiments thereof, as illustrated in the accompanying drawings. The drawings are not necessarily to scale, emphasis instead being placed upon illustrating the principles of the devices, systems, and methods described herein.

(2) FIG. 1 illustrates an environment for threat management.

(3) FIG. 2 illustrates a computer system.

(4) FIG. 3 illustrates a threat management system.

(5) FIG. 4 illustrates a system for behavioral tracking, coloring, and generation of indications of compromise (IOCs).

(6) FIG. 5 illustrates a system for encryption management.

(7) FIG. 6 illustrates a threat management system using heartbeats.

(8) FIG. 7 shows an architecture for endpoint protection in an enterprise network security system.

(9) FIG. 8 illustrates a system for forensic analysis for computer processes.

(10) FIG. 9 is a flowchart of a method for forensic analysis for computer processes.

(11) FIG. 10 illustrates an event graph.

(12) FIG. 11 shows an architecture for instrumenting an endpoint.

(13) FIG. 12 shows a process for accessing encrypted information on an endpoint.

(14) FIG. 13 illustrates malware attempting to access a decrypted key.

(15) FIG. 14 illustrates malware attempting to access a decrypted key.

(16) FIG. 15 illustrates malware attempting to access a decrypted key.

(17) FIG. 16 shows a method for detecting access to an encrypted secret stored on a compute instance.

DETAILED DESCRIPTION

(18) Embodiments will now be described with reference to the accompanying figures, in which preferred embodiments are shown. The foregoing may, however, be embodied in many different forms and should not be construed as limited to the illustrated embodiments set forth herein.

(19) All documents mentioned herein are hereby incorporated by reference in their entirety. References to items in the singular should be understood to include items in the plural, and vice versa, unless explicitly stated otherwise or clear from the context. Grammatical conjunctions are intended to express any and all disjunctive and conjunctive combinations of conjoined clauses, sentences, words, and the like, unless otherwise stated or clear from the context. Thus, the term “or” should generally be understood to mean “and/or” and so forth.

(20) Recitation of ranges of values herein are not intended to be limiting, referring instead individually to any and all values falling within the range, unless otherwise indicated herein, and each separate value within such a range is incorporated into the specification as if it were individually recited herein. The words “about,” “approximately,” or the like, when accompanying a numerical value, are to be construed as indicating a deviation as would be appreciated by one of ordinary skill in the art to operate satisfactorily for an intended purpose. Ranges of values and/or numeric values are provided herein as examples only, and do not constitute a limitation on the scope of the described embodiments. The use of any and all examples, or exemplary language (“e.g.,” “such as,” or the like) provided herein, is intended merely to better illuminate the embodiments and does not pose a limitation on the scope of the embodiments or the claims. No

language in the specification should be construed as indicating any unclaimed element as essential to the practice of the embodiments.

(21) In the following description, it is understood that terms such as “first,” “second,” “third,” “above,” “below,” and the like, are words of convenience and are not to be construed as implying a chronological order or otherwise limiting any corresponding element unless expressly state otherwise.

(22) FIG. 1 illustrates an environment for threat management. Specifically, FIG. 1 depicts a block diagram of a threat management system providing protection to an enterprise against a plurality of threats—a context in which the following techniques may usefully be deployed. One aspect relates to corporate policy management and implementation through a unified threat management facility **100**. As will be explained in more detail below, a threat management facility **100** may be used to protect computer assets from many threats, both computer-generated threats and user-generated threats. The threat management facility **100** may be multi-dimensional in that it may be designed to protect corporate assets from a variety of threats and it may be adapted to learn about threats in one dimension (e.g., worm detection) and apply the knowledge in another dimension (e.g., spam detection). Policy management is one of the dimensions for which the threat management facility can provide a control capability. A corporation or other entity may institute a policy that prevents certain people (e.g., employees, groups of employees, types of employees, guest of the corporation, etc.) from accessing certain types of computer programs. For example, the corporation may elect to prevent its accounting department from using a particular version of an instant messaging service or all such services. In this example, the policy management facility **112** may be used to update the policies of all corporate computing assets with a proper policy control facility or it may update a select few. By using the threat management facility **100** to facilitate the setting, updating and control of such policies the corporation only needs to be concerned with keeping the threat management facility **100** up to date on such policies. The threat management facility **100** can take care of updating all of the other corporate computing assets.

(23) It should be understood that the threat management facility **100** may provide multiple services, and policy management may be offered as one of the services. We will now turn to a description of certain capabilities and components of the threat management system **100**.

(24) Over recent years, malware has become a major problem across the Internet **154**. From both a technical perspective and a user perspective, the categorization of a specific threat type, whether as virus, worm, spam, phishing exploration, spyware, adware, or the like, is becoming reduced in significance. The threat, no matter how it is categorized, may need to be stopped at various points of a networked computing environment, such as one of an enterprise facility **102**, including at one or more laptops, desktops, servers, gateways, communication ports, handheld or mobile devices, firewalls, and the like. Similarly, there may be less and less benefit to the user in having different solutions for known and unknown threats. As such, a consolidated threat management facility **100** may need to apply a similar set of technologies and capabilities for all threats. In certain embodiments, the threat management facility **100** may provide a single agent on the desktop, and a single scan of any suspect file. This approach may eliminate the inevitable overlaps and gaps in protection caused by treating viruses and spyware as separate problems, while simultaneously simplifying administration and minimizing desktop load. As the number and range of types of threats has increased, so may have the level of connectivity available to all IT users. This may have led to a rapid increase in the speed at which threats may move. Today, an unprotected PC connected to the Internet **154** may be infected quickly (perhaps within 10 minutes) which may require acceleration for the delivery of threat protection. Where once monthly updates may have been sufficient, the threat management facility **100** may automatically and seamlessly update its product set against spam and virus threats quickly, for instance, every five minutes, every minute, continuously, or the like. Analysis and testing may be increasingly automated, and also may be performed more frequently; for instance, it may be completed in 15 minutes, and may do so

without compromising quality. The threat management facility **100** may also extend techniques that may have been developed for virus and malware protection and provide them to enterprise facility **102** network administrators to better control their environments. In addition to stopping malicious code, the threat management facility **100** may provide policy management that may be able to control legitimate applications, such as VoIP, instant messaging, peer-to-peer file-sharing, and the like, that may undermine productivity and network performance within the enterprise facility **102**.

(25) The threat management facility **100** may provide an enterprise facility **102** protection from computer-based malware, including viruses, spyware, adware, Trojans, intrusion, spam, policy abuse, uncontrolled access, and the like, where the enterprise facility **102** may be any entity with a networked computer-based infrastructure. In an embodiment, FIG. **1** may depict a block diagram of the threat management facility **100** providing protection to an enterprise against a plurality of threats. The enterprise facility **102** may be corporate, commercial, educational, governmental, or the like, and the enterprise facility's **102** computer network may be distributed amongst a plurality of facilities, and in a plurality of geographical locations, and may include administration **134**, a firewall **138A**, an appliance **140A**, server **142A**, network devices **148A-B**, clients **144A-D**, such as protected by computer security facilities **152**, and the like. It will be understood that any reference herein to client facilities may include the clients **144A-D** shown in FIG. **1** and vice-versa. The threat management facility **100** may include a plurality of functions, such as security management facility **122**, policy management facility **112**, update facility **120**, definitions facility **114**, network access rules facility **124**, remedial action facility **128**, detection techniques facility **130**, testing facility **118**, threat research facility **132**, and the like. In embodiments, the threat protection provided by the threat management facility **100** may extend beyond the network boundaries of the enterprise facility **102** to include clients **144D** (or client facilities) that have moved into network connectivity not directly associated or controlled by the enterprise facility **102**. Threats to client facilities may come from a plurality of sources, such as from network threats **104**, physical proximity threats **110**, secondary location threats **108**, and the like. Clients **144A-D** may be protected from threats even when the client **144A-D** is not located in association with the enterprise **102**, such as when a client **144E-F** moves in and out of the enterprise facility **102**, for example when interfacing with an unprotected server **142C** through the Internet **154**, when a client **144F** is moving into a secondary location threat **108** such as interfacing with components **140B**, **142B**, **148C**, **148D** that are not protected, and the like. In embodiments, the threat management facility **100** may provide an enterprise facility **102** protection from a plurality of threats to multiplatform computer resources in a plurality of locations and network configurations, with an integrated system approach. It should be understood that an enterprise model is applicable to organizations and users of any size or type. For example, an enterprise may be or may include a group or association of endpoints, networks, users, and the like within or outside of one or more protected locations. It should be understood that an enterprise may include one or more offices or business locations, or one or more homes, where each location, or portions of each location, or a collection of locations may be treated as a client facility.

(26) In embodiments, the threat management facility **100** may be provided as a stand-alone solution. In other embodiments, the threat management facility **100** may be integrated into a third-party product. An application programming interface (e.g., a source code interface) may be provided such that the threat management facility **100** may be integrated. For instance, the threat management facility **100** may be stand-alone in that it provides direct threat protection to an enterprise or computer resource, where protection is subscribed to directly **100**. Alternatively, the threat management facility **100** may offer protection indirectly, through a third-party product, where an enterprise may subscribe to services through the third-party product, and threat protection to the enterprise may be provided by the threat management facility **100** through the third-party product.

(27) The security management facility **122** may include a plurality of elements that provide

protection from malware to enterprise facility **102** computer resources, including endpoint security and control, email security and control, web security and control, reputation-based filtering, control of unauthorized users, control of guest and non-compliant computers, and the like. The security management facility **122** may be a software application that may provide malicious code and malicious application protection to a client facility computing resource. The security management facility **122** may have the ability to scan the client facility files for malicious code, remove or quarantine certain applications and files, prevent certain actions, perform remedial actions and perform other security measures. In embodiments, scanning the client facility may include scanning some or all of the files stored to the client facility on a periodic basis, scanning an application when the application is executed, scanning files as the files are transmitted to or from the client facility, or the like. The scanning of the applications and files may be performed to detect known malicious code or known unwanted applications. In an embodiment, new malicious code and unwanted applications may be continually developed and distributed, and updates to the known code database may be provided on a periodic basis, on a demand basis, on an alert basis, or the like.

(28) The security management facility **122** may provide email security and control, where security management may help to eliminate spam, viruses, spyware and phishing, control of email content, and the like. The security management facility's **122** email security and control may protect against inbound and outbound threats, protect email infrastructure, prevent data leakage, provide spam filtering, and the like. In an embodiment, security management facility **122** may provide for web security and control, where security management may help to detect or block viruses, spyware, malware, unwanted applications, help control web browsing, and the like, which may provide comprehensive web access control enabling safe, productive web browsing. Web security and control may provide Internet use policies, reporting on suspect devices, security and content filtering, active monitoring of network traffic, URI filtering, and the like. In an embodiment, the security management facility **122** may provide for network access control, which may provide control over network connections. Network control may stop unauthorized, guest, or non-compliant systems from accessing networks, and may control network traffic that may not be bypassed from the client level. In addition, network access control may control access to virtual private networks (VPN), where VPNs may be a communications network tunneled through another network, establishing a logical connection acting as a virtual network. In embodiments, a VPN may be treated in the same manner as a physical network.

(29) The security management facility **122** may provide host intrusion prevention through behavioral based protection, which may guard against unknown threats by analyzing behavior before software code executes. Behavioral based protection may monitor code when it runs and intervene if the code is deemed to be suspicious or malicious. Advantages of behavioral based protection over runtime protection may include code being prevented from running. Whereas runtime protection may only interrupt code that has already partly executed, behavioral protection can identify malicious code at the gateway or on the file servers and delete the code before it can reach endpoint computers and the like.

(30) The security management facility **122** may provide reputation filtering, which may target or identify sources of known malware. For instance, reputation filtering may include lists of URIs of known sources of malware or known suspicious IP addresses, or domains, say for spam, that when detected may invoke an action by the threat management facility **100**, such as dropping them immediately. By dropping the source before any interaction can initiate, potential threat sources may be thwarted before any exchange of data can be made.

(31) In embodiments, information may be sent from the enterprise back to a third party, a vendor, or the like, which may lead to improved performance of the threat management facility **100**. For example, the types, times, and number of virus interactions that a client experiences may provide useful information for the preventions of future virus threats. This type of feedback may be useful for any aspect of threat detection. Feedback of information may also be associated with behaviors

of individuals within the enterprise, such as being associated with most common violations of policy, network access, unauthorized application loading, unauthorized external device use, and the like. In embodiments, this type of information feedback may enable the evaluation or profiling of client actions that are violations of policy that may provide a predictive model for the improvement of enterprise policies.

(32) The security management facility **122** may support overall security of the enterprise facility **102** network or set of enterprise facility **102** networks, e.g., by providing updates of malicious code information to the enterprise facility **102** network and associated client facilities. The updates may include a planned update, an update in reaction to a threat notice, an update in reaction to a request for an update, an update based on a search of known malicious code information, or the like. The administration facility **134** may provide control over the security management facility **122** when updates are performed. The updates may be automatically transmitted without an administration facility's **134** direct control, manually transmitted by the administration facility **134**, or otherwise distributed. The security management facility **122** may manage the receipt of malicious code descriptions from a provider, distribution of the malicious code descriptions to enterprise facility **102** networks, distribution of the malicious code descriptions to client facilities, and so forth.

(33) The threat management facility **100** may provide a policy management facility **112** that may be able to block non-malicious applications, such as VoIP, instant messaging, peer-to-peer file-sharing, and the like, that may undermine productivity and network performance within the enterprise facility **102**. The policy management facility **112** may be a set of rules or policies that may indicate enterprise facility **102** access permissions for the client facility, such as access permissions associated with the network, applications, external computer devices, and the like. The policy management facility **112** may include a database, a text file, a combination of databases and text files, or the like. In an embodiment, a policy database may be a block list, a black list, an allowed list, a white list, or the like that may provide a list of enterprise facility **102** external network locations/applications that may or may not be accessed by the client facility. The policy management facility **112** may include rules that may be interpreted with respect to an enterprise facility **102** network access request to determine if the request should be allowed. The rules may provide a generic rule for the type of access that may be granted. The rules may be related to the policies of an enterprise facility **102** for access rights for the enterprise facility's **102** client facility. For example, there may be a rule that does not permit access to sporting websites. When a website is requested by the client facility, a security facility may access the rules within a policy facility to determine if the requested access is related to a sporting website. In an embodiment, the security facility may analyze the requested website to determine if the website matches with any of the policy facility rules.

(34) The policy management facility **112** may be similar to the security management facility **122** but with the addition of enterprise facility **102** wide access rules and policies that may be distributed to maintain control of client facility access to enterprise facility **102** network resources. The policies may be defined for application type, subset of application capabilities, organization hierarchy, computer facility type, user type, network location, time of day, connection type, or the like. Policies may be maintained by the administration facility **134**, through the threat management facility **100**, in association with a third party, or the like. For example, a policy may restrict IM activity to only support personnel for communicating with customers. This may allow communication for departments requiring access, but may maintain the network bandwidth for other activities by restricting the use of IM to only the personnel that need access to instant messaging (IM) in support of the enterprise facility **102**. In an embodiment, the policy management facility **112** may be a stand-alone application, may be part of the network server facility **142**, may be part of the enterprise facility **102** network, may be part of the client facility, or the like.

(35) The threat management facility **100** may provide configuration management, which may be similar to policy management, but may specifically examine the configuration set of applications,

operating systems, hardware, and the like, and manage changes to their configurations. Assessment of a configuration may be made against a standard configuration policy, detection of configuration changes, remediation of improper configuration, application of new configurations, and the like. An enterprise may keep a set of standard configuration rules and policies which may represent the desired state of the device. For example, a client firewall may be running and installed, but in the disabled state, where remediation may be to enable the firewall. In another example, the enterprise may set a rule that disallows the use of USB disks, and sends a configuration change to all clients, which turns off USB drive access via a registry.

(36) The threat management facility **100** may also provide for the removal of applications that potentially interfere with the operation of the threat management facility **100**, such as competitor products that may also be attempting similar threat management functions. The removal of such products may be initiated automatically whenever such products are detected. In the case where such applications are services are provided indirectly through a third-party product, the application may be suspended until action is taken to remove or disable the third-party product's protection facility.

(37) Threat management against a quickly evolving malware environment may require timely updates, and thus an update management facility **120** may be provided by the threat management facility **100**. In addition, a policy management facility **112** may also require update management (e.g., as provided by the update facility **120** herein described). The update management for the security facility **122** and policy management facility **112** may be provided directly by the threat management facility **100**, such as by a hosted system or in conjunction with the administration facility **134**. In embodiments, the threat management facility **100** may provide for patch management, where a patch may be an update to an operating system, an application, a system tool, or the like, where one of the reasons for the patch is to reduce vulnerability to threats.

(38) The security facility **122** and policy management facility **112** may push information to the enterprise facility **102** network and/or client facility. The enterprise facility **102** network and/or client facility may also or instead pull information from the security facility **122** and policy management facility **112** network server facilities **142**, or there may be a combination of pushing and pulling of information between the security facility **122** and the policy management facility **112** network servers **142**, enterprise facility **102** network, and client facilities, or the like. For example, the enterprise facility **102** network and/or client facility may pull information from the security facility **122** and policy management facility **112** network server facility **142** may request the information using the security facility **122** and policy management facility **112** update module; the request may be based on a certain time period, by a certain time, by a date, on demand, or the like. In another example, the security facility **122** and policy management facility **112** network servers **142** may push the information to the enterprise facility's **102** network and/or client facility by providing notification that there are updates available for download and then transmitting the information. The combination of the security management **122** network server facility **142** and security update module may function substantially the same as the policy management facility **112** network server and policy update module by providing information to the enterprise facility **102** network and the client facility in a push or pull method. In an embodiment, the policy management facility **112** and the security facility **122** management update modules may work in concert to provide information to the enterprise facility's **102** network and/or client facility for control of application execution. In an embodiment, the policy update module and security update module may be combined into a single update module.

(39) As threats are identified and characterized, the threat management facility **100** may create definition updates that may be used to allow the threat management facility **100** to detect and remediate the latest malicious software, unwanted applications, configuration and policy changes, and the like. The threat definition facility **114** may contain threat identification updates, also referred to as definition files. A definition file may be a virus identity file that may include

definitions of known or potential malicious code. The virus identity (IDE) definition files may provide information that may identify malicious code within files, applications, or the like. The definition files may be accessed by security management facility **122** when scanning files or applications within the client facility for the determination of malicious code that may be within the file or application. The definition files may contain a number of commands, definitions, or instructions, to be parsed and acted upon, or the like. In embodiments, the client facility may be updated with new definition files periodically to provide the client facility with the most recent malicious code definitions; the updating may be performed on a set time period, may be updated on demand from the client facility, may be updated on demand from the network, may be updated on a received malicious code alert, or the like. In an embodiment, the client facility may request an update to the definition files from an update facility **120** within the network, may request updated definition files from a computing facility external to the network, updated definition files may be provided to the client facility **114** from within the network, definition files may be provided to the client facility from an external computing facility from an external network, or the like.

(40) A definition management facility **114** may provide timely updates of definition files information to the network, client facilities, and the like. New and altered malicious code and malicious applications may be continually created and distributed to networks worldwide. The definition files that maintain the definitions of the malicious code and malicious application information for the protection of the networks and client facilities may need continual updating to provide continual defense of the network and client facility from the malicious code and malicious applications. The definition files management may provide for automatic and manual methods of updating the definition files. In embodiments, the network may receive definition files and distribute the definition files to the network client facilities, the client facilities may receive the definition files directly, or the network and client facilities may both receive the definition files, or the like. In an embodiment, the definition files may be updated on a fixed periodic basis, on demand by the network and/or the client facility, as a result of an alert of a new malicious code or malicious application, or the like. In an embodiment, the definition files may be released as a supplemental file to an existing definition files to provide for rapid updating of the definition files.

(41) In a similar manner, the security management facility **122** may be used to scan an outgoing file and verify that the outgoing file is permitted to be transmitted per the enterprise facility **102** rules and policies. By checking outgoing files, the security management facility **122** may be able to discover malicious code infected files that were not detected as incoming files as a result of the client facility having been updated with either new definition files or policy management facility **112** information. The definition files may discover the malicious code infected file by having received updates of developing malicious code from the administration facility **134**, updates from a definition files provider, or the like. The policy management facility **112** may discover the malicious code infected file by having received new updates from the administration facility **134**, from a rules provider, or the like.

(42) The threat management facility **100** may provide controlled access to the enterprise facility **102** networks. For instance, a manager of the enterprise facility **102** may want to restrict access to certain applications, networks, files, printers, servers, databases, or the like. In addition, the manager of the enterprise facility **102** may want to restrict user access based on certain criteria, such as the user's location, usage history, need to know, job position, connection type, time of day, method of authentication, client-system configuration, or the like. Network access rules may be developed for the enterprise facility **102**, or pre-packaged by a supplier, and managed by the threat management facility **100** in conjunction with the administration facility **134**.

(43) A network access rules facility **124** may be responsible for determining if a client facility application should be granted access to a requested network location. The network location may be on the same network as the facility or may be on another network. In an embodiment, the network access rules facility **124** may verify access rights for client facilities from within the network or

may verify access rights of computer facilities from external networks. When network access for a client facility is denied, the network access rules facility **124** may send an information file to the client facility containing. For example, the information sent by the network access rules facility **124** may be a data file. The data file may contain a number of commands, definitions, instructions, or the like to be parsed and acted upon through the remedial action facility **128**, or the like. The information sent by the network access facility rules facility **124** may be a command or command file that the remedial action facility **128** may access and take action upon.

(44) The network access rules facility **124** may include databases such as a block list, a black list, an allowed list, a white list, an unacceptable network site database, an acceptable network site database, a network site reputation database, or the like of network access locations that may or may not be accessed by the client facility. Additionally, the network access rules facility **124** may incorporate rule evaluation; the rule evaluation may parse network access requests and apply the parsed information to network access rules. The network access rule facility **124** may have a generic set of rules that may be in support of an enterprise facility's **102** network access policies, such as denying access to certain types of websites, controlling instant messenger accesses, or the like. Rule evaluation may include regular expression rule evaluation, or other rule evaluation method for interpreting the network access request and comparing the interpretation to the established rules for network access. In an embodiment, the network access rules facility **124** may receive a rules evaluation request from the network access control and may return the rules evaluation to the network access control.

(45) Similar to the threat definitions facility **114**, the network access rule facility **124** may provide updated rules and policies to the enterprise facility **102**. The network access rules facility **124** may be maintained by the network administration facility **134**, using network access rules facility **124** management. In an embodiment, the network administration facility **134** may be able to maintain a set of access rules manually by adding rules, changing rules, deleting rules, or the like. Additionally, the administration facility **134** may retrieve predefined rule sets from a remote provider of a set of rules to be applied to an entire enterprise facility **102**. The network administration facility **134** may be able to modify the predefined rules as needed for a particular enterprise facility **102** using the network access rules management facility **124**.

(46) When a threat or policy violation is detected by the threat management facility **100**, the threat management facility **100** may perform or initiate a remedial action facility **128**. Remedial action may take a plurality of forms, such as terminating or modifying an ongoing process or interaction, sending a warning to a client or administration facility **134** of an ongoing process or interaction, executing a program or application to remediate against a threat or violation, record interactions for subsequent evaluation, or the like. Remedial action may be associated with an application that responds to information that a client facility network access request has been denied. In an embodiment, when the data file is received, remedial action may parse the data file, interpret the various aspects of the data file, and act on the parsed data file information to determine actions to be taken on an application requesting access to a denied network location. In an embodiment, when the data file is received, remedial action may access the threat definitions to parse the data file and determine an action to be taken on an application requesting access to a denied network location. In an embodiment, the information received from the facility may be a command or a command file. The remedial action facility may carry out any commands that are received or parsed from a data file from the facility without performing any interpretation of the commands. In an embodiment, the remedial action facility may interact with the received information and may perform various actions on a client requesting access to a denied network location. The action may be one or more of continuing to block all requests to a denied network location, a malicious code scan on the application, a malicious code scan on the client facility, quarantine of the application, terminating the application, isolation of the application, isolation of the client facility to a location within the network that restricts network access, blocking a network access port from a client facility,

reporting the application to an administration facility **134**, or the like.

(47) Remedial action may be provided as a result of a detection of a threat or violation. The detection techniques facility **130** may include monitoring the enterprise facility **102** network or endpoint devices, such as by monitoring streaming data through the gateway, across the network, through routers and hubs, and the like. The detection techniques facility **130** may include monitoring activity and stored files on computing facilities, such as on server facilities **142**, desktop computers, laptop computers, other mobile computing devices, and the like. Detection techniques, such as scanning a computer's stored files, may provide the capability of checking files for stored threats, either in the active or passive state. Detection techniques, such as streaming file management, may provide the capability of checking files received at the network, gateway facility, client facility, and the like. This may provide the capability of not allowing a streaming file or portions of the streaming file containing malicious code from entering the client facility, gateway facility, or network. In an embodiment, the streaming file may be broken into blocks of information, and a plurality of virus identities may be used to check each of the blocks of information for malicious code. In an embodiment, any blocks that are not determined to be clear of malicious code may not be delivered to the client facility, gateway facility, or network.

(48) Verifying that the threat management facility **100** is detecting threats and violations to established policy, may require the ability to test the system, either at the system level or for a particular computing component. The testing facility **118** may allow the administration facility **134** to coordinate the testing of the security configurations of client facility computing facilities on a network. The administration facility **134** may be able to send test files to a set of client facility computing facilities to test the ability of the client facility to determine acceptability of the test file. After the test file has been transmitted, a recording facility may record the actions taken by the client facility in reaction to the test file. The recording facility may aggregate the testing information from the client facility and report the testing information to the administration facility **134**. The administration facility **134** may be able to determine the level of preparedness of the client facility computing facilities by the reported information. Remedial action may be taken for any of the client facility computing facilities as determined by the administration facility **134**; remedial action may be taken by the administration facility **134** or by the user of the client facility.

(49) The threat research facility **132** may provide a continuously ongoing effort to maintain the threat protection capabilities of the threat management facility **100** in light of continuous generation of new or evolved forms of malware. Threat research may include researchers and analysts working on known and emerging malware, such as viruses, rootkits a spyware, as well as other computer threats such as phishing, spam, scams, and the like. In embodiments, through threat research, the threat management facility **100** may be able to provide swift, global responses to the latest threats.

(50) The threat management facility **100** may provide threat protection to the enterprise facility **102**, where the enterprise facility **102** may include a plurality of networked components, such as client facility, server facility **142**, administration facility **134**, firewall **138**, gateway, hubs, and routers **148**, threat management appliance **140**, desktop users, mobile users, and the like. In embodiments, it may be the endpoint computer security facility **152**, located on a computer's desktop, which may provide threat protection to a user, and associated enterprise facility **102**. In embodiments, the term endpoint may refer to a computer system that may source data, receive data, evaluate data, buffer data, or the like (such as a user's desktop computer as an endpoint computer), a firewall as a data evaluation endpoint computer system, a laptop as a mobile endpoint computer, a personal digital assistant or tablet as a hand-held endpoint computer, a mobile phone as an endpoint computer, or the like. In embodiments, endpoint may refer to a source or destination for data, including such components where the destination is characterized by an evaluation point for data, and where the data may be sent to a subsequent destination after evaluation. The endpoint computer security facility **152** may be an application loaded onto the computer platform or computer support

component, where the application may accommodate the plurality of computer platforms and/or functional requirements of the component. For instance, a client facility computer may be one of a plurality of computer platforms, such as Windows, Macintosh, Linux, and the like, where the endpoint computer security facility **152** may be adapted to the specific platform, while maintaining a uniform product and product services across platforms. Additionally, components may have different functions to serve within the enterprise facility's **102** networked computer-based infrastructure. For instance, computer support components provided as hubs and routers **148**, server facility **142**, firewalls **138**, and the like, may require unique security application software to protect their portion of the system infrastructure, while providing an element in an integrated threat management system that extends out beyond the threat management facility **100** to incorporate all computer resources under its protection.

(51) The enterprise facility **102** may include a plurality of client facility computing platforms on which the endpoint computer security facility **152** is adapted. A client facility computing platform may be a computer system that is able to access a service on another computer, such as a server facility **142**, via a network. This client facility server facility **142** model may apply to a plurality of networked applications, such as a client facility connecting to an enterprise facility **102** application server facility **142**, a web browser client facility connecting to a web server facility **142**, an e-mail client facility retrieving e-mail from an Internet **154** service provider's mail storage servers **142**, and the like. In embodiments, traditional large client facility applications may be switched to websites, which may increase the browser's role as a client facility. Clients **144** may be classified as a function of the extent to which they perform their own processing. For instance, client facilities are sometimes classified as a fat client facility or thin client facility. The fat client facility, also known as a thick client facility or rich client facility, may be a client facility that performs the bulk of data processing operations itself, and does not necessarily rely on the server facility **142**. The fat client facility may be most common in the form of a personal computer, where the personal computer may operate independent of any server facility **142**. Programming environments for fat clients **144** may include CURI, Delphi, Droplets, Java, win32, X11, and the like. Thin clients **144** may offer minimal processing capabilities, for instance, the thin client facility may primarily provide a graphical user interface provided by an application server facility **142**, which may perform the bulk of any required data processing. Programming environments for thin clients **144** may include JavaScript/AJAX, ASP, JSP, Ruby on Rails, Python's Django, PHP, and the like. The client facility may also be a mix of the two, such as processing data locally, but relying on a server facility **142** for data storage. As a result, this hybrid client facility may provide benefits from both the fat client facility type, such as multimedia support and high performance, and the thin client facility type, such as high manageability and flexibility. In embodiments, the threat management facility **100**, and associated endpoint computer security facility **152**, may provide seamless threat protection to the plurality of clients **144**, and client facility types, across the enterprise facility **102**.

(52) The enterprise facility **102** may include a plurality of server facilities **142**, such as application servers, communications servers, file servers, database servers, proxy servers, mail servers, fax servers, game servers, web servers, and the like. A server facility **142**, which may also be referred to as a server facility **142** application, server facility **142** operating system, server facility **142** computer, or the like, may be an application program or operating system that accepts client facility connections in order to service requests from clients **144**. The server facility **142** application may run on the same computer as the client facility using it, or the server facility **142** and the client facility may be running on different computers and communicating across the network. Server facility **142** applications may be divided among server facility **142** computers, with the dividing depending upon the workload. For instance, under light load conditions all server facility **142** applications may run on a single computer and under heavy load conditions a single server facility **142** application may run on multiple computers. In embodiments, the threat management facility **100** may provide threat protection to server facilities **142** within the enterprise facility **102** as load

conditions and application changes are made.

(53) A server facility **142** may also be an appliance facility **140**, where the appliance facility **140** provides specific services onto the network. Though the appliance facility **140** is a server facility **142** computer, that may be loaded with a server facility **142** operating system and server facility **142** application, the enterprise facility **102** user may not need to configure it, as the configuration may have been performed by a third party. In an embodiment, an enterprise facility **102** appliance may be a server facility **142** appliance that has been configured and adapted for use with the threat management facility **100**, and located within the facilities of the enterprise facility **102**. The enterprise facility's **102** threat management appliance may enable the enterprise facility **102** to administer an on-site local managed threat protection configuration, where the administration facility **134** may access the threat resources through an interface, such as a web portal. In an alternate embodiment, the enterprise facility **102** may be managed remotely from a third party, vendor, or the like, without an appliance facility **140** located within the enterprise facility **102**. In this instance, the appliance functionality may be a shared hardware product between pluralities of enterprises **102**. In embodiments, the appliance facility **140** may be located at the enterprise facility **102**, where the enterprise facility **102** maintains a degree of control. In embodiments, a hosted service may be provided, where the appliance **140** may still be an on-site black box to the enterprise facility **102**, physically placed there because of infrastructure requirements, but managed by a third party, vendor, or the like.

(54) Simple server facility **142** appliances may also be utilized across the enterprise facility's **102** network infrastructure, such as switches, routers, wireless routers, hubs and routers, gateways, print servers, net modems, and the like. These simple server facility appliances may not require configuration by the enterprise facility **102**, but may require protection from threats via an endpoint computer security facility **152**. These appliances may provide interconnection services within the enterprise facility **102** network, and therefore may advance the spread of a threat if not properly protected.

(55) A client facility may be protected from threats from within the enterprise facility **102** network using a personal firewall, which may be a hardware firewall, software firewall, or combination of these, that controls network traffic to and from a client. The personal firewall may permit or deny communications based on a security policy. Personal firewalls may be designed for use by end-users, which may result in protection for only the computer on which it's installed. Personal firewalls may be able to control network traffic by providing prompts each time a connection is attempted and adapting security policy accordingly. Personal firewalls may also provide some level of intrusion detection, which may allow the software to terminate or block connectivity where it suspects an intrusion is being attempted. Other features that may be provided by a personal firewall may include alerts about outgoing connection attempts, control of program access to networks, hiding the client from port scans by not responding to unsolicited network traffic, monitoring of applications that may be listening for incoming connections, monitoring and regulation of incoming and outgoing network traffic, prevention of unwanted network traffic from installed applications, reporting applications that make connection attempts, reporting destination servers with which applications may be attempting communications, and the like. In embodiments, the personal firewall may be provided by the threat management facility **100**.

(56) Another important component that may be protected by an endpoint computer security facility **152** is a network firewall facility **138**, which may be a hardware or software device that may be configured to permit, deny, or proxy data through a computer network that has different levels of trust in its source of data. For instance, an internal enterprise facility **102** network may have a high level of trust, because the source of all data has been sourced from within the enterprise facility **102**. An example of a low level of trust is the Internet **154**, because the source of data may be unknown. A zone with an intermediate trust level, situated between the Internet **154** and a trusted internal network, may be referred to as a "perimeter network." Since firewall facilities **138**

represent boundaries between threat levels, the endpoint computer security facility **152** associated with the firewall facility **138** may provide resources that may control the flow of threats at this enterprise facility **102** network entry point. Firewall facilities **138**, and associated endpoint computer security facility **152**, may also be associated with a network node that may be equipped for interfacing between networks that use different protocols. In embodiments, the endpoint computer security facility **152** may provide threat protection in a plurality of network infrastructure locations, such as at the enterprise facility **102** network entry point, i.e., the firewall facility **138** or gateway; at the server facility **142**; at distribution points within the network, i.e., the hubs and routers **148**; at the desktop of client facility computers; and the like. In embodiments, the most effective location for threat detection may be at the user's computer desktop endpoint computer security facility **152**.

(57) The interface between the threat management facility **100** and the enterprise facility **102**, and through the appliance facility **140** to embedded endpoint computer security facilities, may include a set of tools that may be the same for all enterprise implementations, but allow each enterprise to implement different controls. In embodiments, these controls may include both automatic actions and managed actions. Automatic actions may include downloads of the endpoint computer security facility **152** to components of the enterprise facility **102**, downloads of updates to existing endpoint computer security facilities of the enterprise facility **102**, uploaded network interaction requests from enterprise facility **102** components to the threat management facility **100**, and the like. In embodiments, automatic interactions between the enterprise facility **102** and the threat management facility **100** may be configured by the threat management facility **100** and an administration facility **134** in the enterprise facility **102**. The administration facility **134** may configure policy rules that determine interactions, such as developing rules for accessing applications, as in who is authorized and when applications may be used; establishing rules for ethical behavior and activities; rules governing the use of entertainment software such as games, or personal use software such as IM and VoIP; rules for determining access to enterprise facility **102** computing resources, including authentication, levels of access, risk assessment, and usage history tracking; rules for when an action is not allowed, such as whether an action is completely deigned or just modified in its execution; and the like. The administration facility **134** may also establish license management, which in turn may further determine interactions associated with a licensed application. In embodiments, interactions between the threat management facility **100** and the enterprise facility **102** may provide threat protection to the enterprise facility **102** by managing the flow of network data into and out of the enterprise facility **102** through automatic actions that may be configured by the threat management facility **100** or the administration facility **134**.

(58) Client facilities within the enterprise facility **102** may be connected to the enterprise facility **102** network by way of wired network facilities **148A** or wireless network facilities **148B**. Client facilities connected to the enterprise facility **102** network via a wired facility **148A** or wireless facility **148B** may receive similar protection, as both connection types are ultimately connected to the same enterprise facility **102** network, with the same endpoint computer security facility **152**, and the same threat protected enterprise facility **102** environment. Mobile wireless facility clients **144B-F**, because of their ability to connect to any wireless **148B**, D network access point, may connect to the Internet **154** outside the enterprise facility **102**, and therefore outside the threat-protected environment of the enterprise facility **102**. In this instance the mobile client facility (e.g., the clients **144B-F**), if not for the presence of the endpoint computer security facility **152** may experience a malware attack or perform actions counter to enterprise facility **102** established policies. In addition, there may be a plurality of ways for the threat management facility **100** to protect the out-of-enterprise facility **102** mobile client facility (e.g., the clients **144D-F**) that has an embedded endpoint computer security facility **152**, such as by providing URI filtering in personal routers, using a web appliance as a DNS proxy, or the like. Mobile client facilities that are components of the enterprise facility **102** but temporarily outside connectivity with the enterprise

facility **102** network may be provided with the same threat protection and policy control as client facilities inside the enterprise facility **102**. In addition, mobile the client facilities may receive the same interactions to and from the threat management facility **100** as client facilities inside the enterprise facility **102**, where the mobile client facilities may be considered a virtual extension of the enterprise facility **102**, receiving all the same services via their embedded endpoint computer security facility **152**.

(59) Interactions between the threat management facility **100** and the components of the enterprise facility **102**, including mobile client facility extensions of the enterprise facility **102**, may ultimately be connected through the Internet **154**. Threat management facility **100** downloads and upgrades to the enterprise facility **102** may be passed from the firewalled networks of the threat management facility **100** through to the endpoint computer security facility **152** equipped components of the enterprise facility **102**. In turn the endpoint computer security facility **152** components of the enterprise facility **102** may upload policy and access requests back across the Internet **154** and through to the threat management facility **100**. The Internet **154** however, is also the path through which threats may be transmitted from their source. These network threats **104** may include threats from a plurality of sources, including without limitation, websites, e-mail, IM, VOIP, application software, and the like. These threats may attempt to attack a mobile enterprise client facility (e.g., the clients **144B-F**) equipped with an endpoint computer security facility **152**, but in embodiments, as long as the mobile client facility is embedded with an endpoint computer security facility **152**, as described above, threats may have no better success than if the mobile client facility were inside the enterprise facility **102**.

(60) However, if the mobile client facility were to attempt to connect into an unprotected connection point, such as at a secondary location **108** that is not a part of the enterprise facility **102**, the mobile client facility may be required to request network interactions through the threat management facility **100**, where contacting the threat management facility **100** may be performed prior to any other network action. In embodiments, the client facility's **144** endpoint computer security facility **152** may manage actions in unprotected network environments such as when the client facility (e.g., client **144F**) is in a secondary location **108** or connecting wirelessly to a non-enterprise facility **102** wireless Internet connection, where the endpoint computer security facility **152** may dictate what actions are allowed, blocked, modified, or the like. For instance, if the client facility's **144** endpoint computer security facility **152** is unable to establish a secured connection to the threat management facility **100**, the endpoint computer security facility **152** may inform the user of such and recommend that the connection not be made. In the instance when the user chooses to connect despite the recommendation, the endpoint computer security facility **152** may perform specific actions during or after the unprotected connection is made, including running scans during the connection period, running scans after the connection is terminated, storing interactions for subsequent threat and policy evaluation, contacting the threat management facility **100** upon first instance of a secured connection for further actions and or scanning, restricting access to network and local resources, or the like. In embodiments, the endpoint computer security facility **152** may perform specific actions to remediate possible threat incursions or policy violations during or after the unprotected connection.

(61) The secondary location **108** may have no endpoint computer security facilities **152** as a part of its computer components, such as its firewalls **138B**, servers **142B**, clients **144G**, hubs and routers **148C-D**, and the like. As a result, the computer components of the secondary location **108** may be open to threat attacks, and become potential sources of threats, as well as any mobile enterprise facility clients **144B-F** that may be connected to the secondary location's **108** network. In this instance, these computer components may now unknowingly spread a threat to other components connected to the network.

(62) Some threats may not come directly from the Internet **154**, such as from non-enterprise facility controlled mobile devices that are physically brought into the enterprise facility **102** and connected

to the enterprise facility **102** client facilities. The connection may be made from direct connection with the enterprise facility's **102** client facility, such as through a USB port, or in physical proximity with the enterprise facility's **102** client facility such that a wireless facility connection can be established, such as through a Bluetooth connection. These physical proximity threats **110** may be another mobile computing device, a portable memory storage device, a mobile communications device, or the like, such as CDs and DVDs, memory sticks, flash drives, external hard drives, cell phones, PDAs, MP3 players, digital cameras, point-to-point devices, digital picture frames, digital pens, navigation devices, tablets, appliances, and the like. A physical proximity threat **110** may have been previously infiltrated by network threats while connected to an unprotected network connection outside the enterprise facility **102**, and when connected to the enterprise facility **102** client facility, pose a threat. Because of their mobile nature, physical proximity threats **110** may infiltrate computing resources in any location, such as being physically brought into the enterprise facility **102** site, connected to an enterprise facility **102** client facility while that client facility is mobile, plugged into an unprotected client facility at a secondary location **108**, and the like. A mobile device, once connected to an unprotected computer resource, may become a physical proximity threat **110**. In embodiments, the endpoint computer security facility **152** may provide enterprise facility **102** computing resources with threat protection against physical proximity threats **110**, for instance, through scanning the device prior to allowing data transfers, through security validation certificates, through establishing a safe zone within the enterprise facility **102** computing resource to transfer data into for evaluation, and the like.

(63) Having provided an overall context for threat detection, the description now turns to a brief discussion of an example of a computer system that may be used for any of the entities and facilities described above.

(64) FIG. 2 illustrates a computer system. In general, the computer system **200** may include a computing device **210** connected to a network **202**, e.g., through an external device **204**. The computing device **210** may be or include any type of network endpoint or endpoints as described herein, e.g., with reference to FIG. 1 above. For example, the computing device **210** may include a desktop computer workstation. The computing device **210** may also or instead be any suitable device that has processes and communicates over a network **202**, including without limitation a laptop computer, a desktop computer, a personal digital assistant, a tablet, a mobile phone, a television, a set top box, a wearable computer (e.g., watch, jewelry, or clothing), a home device (e.g., a thermostat or a home appliance controller), just as some examples. The computing device **210** may also or instead include a server, or it may be disposed on a server.

(65) The computing device **210** may be used for any of the entities described in the threat management environment described above with reference to FIG. 1. For example, the computing device **210** may be a server, a client an enterprise facility, a threat management facility, or any of the other facilities or computing devices described therein. In certain aspects, the computing device **210** may be implemented using hardware (e.g., in a desktop computer), software (e.g., in a virtual machine or the like), or a combination of software and hardware (e.g., with programs executing on the desktop computer), and the computing device **210** may be a standalone device, a device integrated into another entity or device, a platform distributed across multiple entities, or a virtualized device executing in a virtualization environment.

(66) The network **202** may include any network described above, e.g., data network(s) or internetwork(s) suitable for communicating data and control information among participants in the computer system **200**. This may include public networks such as the Internet, private networks, and telecommunications networks such as the Public Switched Telephone Network or cellular networks using third generation cellular technology (e.g., 3G or IMT-2000), fourth generation cellular technology (e.g., 4G, LTE. MT-Advanced, E-UTRA, etc.) or WiMax-Advanced (IEEE 802.16m)) and/or other technologies, as well as any of a variety of corporate area, metropolitan area, campus or other local area networks or enterprise networks, along with any switches, routers, hubs,

gateways, and the like that might be used to carry data among participants in the computer system **200**. The network **202** may also include a combination of data networks and need not be limited to a strictly public or private network.

(67) The external device **204** may be any computer or other remote resource that connects to the computing device **210** through the network **202**. This may include threat management resources such as any of those contemplated above, gateways or other network devices, remote servers or the like containing content requested by the computing device **210**, a network storage device or resource, a device hosting malicious content, or any other resource or device that might connect to the computing device **210** through the network **202**.

(68) The computing device **210** may include a processor **212**, a memory **214**, a network interface **216**, a data store **218**, and one or more input/output devices **220**. The computing device **210** may further include or be in communication with peripherals **222** and other external input/output devices **224**.

(69) The processor **212** may be any as described herein, and in general be capable of processing instructions for execution within the computing device **210** or computer system **200**. The processor **212** may include a single-threaded processor or a multi-threaded processor. The processor **212** may be capable of processing instructions stored in the memory **214** or on the data store **218**.

(70) The memory **214** may store information within the computing device **210** or computer system **200**. The memory **214** may include any volatile or non-volatile memory or other computer-readable medium, including without limitation a Random-Access Memory (RAM), a flash memory, a Read Only Memory (ROM), a Programmable Read-only Memory (PROM), an Erasable PROM (EPROM), registers, and so forth. The memory **214** may store program instructions, program data, executables, and other software and data useful for controlling operation of the computing device **200** and configuring the computing device **200** to perform functions for a user. The memory **214** may include a number of different stages and types for different aspects of operation of the computing device **210**. For example, a processor may include on-board memory and/or cache for faster access to certain data or instructions, and a separate, main memory or the like may be included to expand memory capacity as desired.

(71) The memory **214** may, in general, include a non-volatile computer readable medium containing computer code that, when executed by the computing device **200** creates an execution environment for a computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, or a combination of the foregoing, and/or code that performs some or all of the steps set forth in the various flow charts and other algorithmic descriptions set forth herein. While a single memory **214** is depicted, it will be understood that any number of memories may be usefully incorporated into the computing device **210**. For example, a first memory may provide non-volatile storage such as a disk drive for permanent or long-term storage of files and code even when the computing device **210** is powered down. A second memory such as a random-access memory may provide volatile (but higher speed) memory for storing instructions and data for executing processes. A third memory may be used to improve performance by providing even higher speed memory physically adjacent to the processor **212** for registers, caching and so forth.

(72) The network interface **216** may include any hardware and/or software for connecting the computing device **210** in a communicating relationship with other resources through the network **202**. This may include remote resources accessible through the Internet, as well as local resources available using short range communications protocols using, e.g., physical connections (e.g., Ethernet), radio frequency communications (e.g., WiFi), optical communications, (e.g., fiber optics, infrared, or the like), ultrasonic communications, or any combination of these or other media that might be used to carry data between the computing device **210** and other devices. The network interface **216** may, for example, include a router, a modem, a network card, an infrared transceiver, a radio frequency (RF) transceiver, a near field communications interface, a radio-frequency

identification (RFID) tag reader, or any other data reading or writing resource or the like.

(73) More generally, the network interface **216** may include any combination of hardware and software suitable for coupling the components of the computing device **210** to other computing or communications resources. By way of example and not limitation, this may include electronics for a wired or wireless Ethernet connection operating according to the IEEE 802.11 standard (or any variation thereof), or any other short or long range wireless networking components or the like. This may include hardware for short range data communications such as Bluetooth or an infrared transceiver, which may be used to couple to other local devices, or to connect to a local area network or the like that is in turn coupled to a data network **202** such as the Internet. This may also or instead include hardware/software for a WiMax connection or a cellular network connection (using, e.g., CDMA, GSM, LTE, or any other suitable protocol or combination of protocols). The network interface **216** may be included as part of the input/output devices **220** or vice-versa.

(74) The data store **218** may be any internal memory store providing a computer-readable medium such as a disk drive, an optical drive, a magnetic drive, a flash drive, or other device capable of providing mass storage for the computing device **210**. The data store **218** may store computer readable instructions, data structures, program modules, and other data for the computing device **210** or computer system **200** in a non-volatile form for subsequent retrieval and use. For example, the data store **218** may store without limitation one or more of the operating system, application programs, program data, databases, files, and other program modules or other software objects and the like.

(75) The input/output interface **220** may support input from and output to other devices that might couple to the computing device **210**. This may, for example, include serial ports (e.g., RS-232 ports), universal serial bus (USB) ports, optical ports, Ethernet ports, telephone ports, audio jacks, component audio/video inputs, HDMI ports, and so forth, any of which might be used to form wired connections to other local devices. This may also or instead include an infrared interface, RF interface, magnetic card reader, or other input/output system for coupling in a communicating relationship with other local devices. It will be understood that, while the network interface **216** for network communications is described separately from the input/output interface **220** for local device communications, these two interfaces may be the same, or may share functionality, such as where a USB port is used to attach to a WiFi accessory, or where an Ethernet connection is used to couple to a local network attached storage.

(76) A peripheral **222** may include any device used to provide information to or receive information from the computing device **200**. This may include human input/output (I/O) devices such as a keyboard, a mouse, a mouse pad, a track ball, a joystick, a microphone, a foot pedal, a camera, a touch screen, a scanner, or other device that might be employed by the user **230** to provide input to the computing device **210**. This may also or instead include a display, a speaker, a printer, a projector, a headset, or any other audiovisual device for presenting information to a user. The peripheral **222** may also or instead include a digital signal processing device, an actuator, or other device to support control or communication to other devices or components. Other I/O devices suitable for use as a peripheral **222** include haptic devices, three-dimensional rendering systems, augmented-reality displays, magnetic card readers, and so forth. In one aspect, the peripheral **222** may serve as the network interface **216**, such as with a USB device configured to provide communications via short range (e.g., Bluetooth, WiFi, Infrared, RF, or the like) or long range (e.g., cellular data or WiMax) communications protocols. In another aspect, the peripheral **222** may provide a device to augment operation of the computing device **210**, such as a global positioning system (GPS) device, a security dongle, or the like. In another aspect, the peripheral may be a storage device such as a flash card, USB drive, or other solid-state device, or an optical drive, a magnetic drive, a disk drive, or other device or combination of devices suitable for bulk storage. More generally, any device or combination of devices suitable for use with the computing device **200** may be used as a peripheral **222** as contemplated herein.

(77) Other hardware **226** may be incorporated into the computing device **200** such as a co-processor, a digital signal processing system, a math co-processor, a graphics engine, a video driver, and so forth. The other hardware **226** may also or instead include expanded input/output ports, extra memory, additional drives (e.g., a DVD drive or other accessory), and so forth.

(78) A bus **232** or combination of busses may serve as an electromechanical platform for interconnecting components of the computing device **200** such as the processor **212**, memory **214**, network interface **216**, other hardware **226**, data store **218**, and input/output interface. As shown in the figure, each of the components of the computing device **210** may be interconnected using a system bus **232** or other communication mechanism for communicating information.

(79) Methods and systems described herein can be realized using the processor **212** of the computer system **200** to execute one or more sequences of instructions contained in the memory **214** to perform predetermined tasks. In embodiments, the computing device **200** may be deployed as a number of parallel processors synchronized to execute code together for improved performance, or the computing device **200** may be realized in a virtualized environment where software on a hypervisor or other virtualization management facility emulates components of the computing device **200** as appropriate to reproduce some or all of the functions of a hardware instantiation of the computing device **200**.

(80) FIG. 3 illustrates a threat management system according to some implementations. In general, the system **300** may include an endpoint **302**, a firewall **304**, a server **306** and a threat management facility **308** coupled to one another directly or indirectly through a data network **305**, all as generally described above. Each of the entities depicted in FIG. 3 may, for example, be implemented on one or more computing devices such as the computing device described above with reference to FIG. 2. A number of systems may be distributed across these various components to support threat detection, such as a coloring system **310**, a key management system **312** and a heartbeat system **314** (or otherwise an endpoint health system), each of which may include software components executing on any of the foregoing system components, and each of which may communicate with the threat management facility **308** and an endpoint threat detection agent **320** executing on the endpoint **302** to support improved threat detection and remediation.

(81) The coloring system **310** may be used to label or ‘color’ software objects for improved tracking and detection of potentially harmful activity. The coloring system **310** may, for example, label files, executables, processes, network communications, data sources and so forth with any suitable label. A variety of techniques may be used to select static and/or dynamic labels for any of these various software objects, and to manage the mechanics of applying and propagating coloring information as appropriate. For example, a process may inherit a color from an application that launches the process. Similarly, a file may inherit a color from a process when it is created or opened by a process, and/or a process may inherit a color from a file that the process has opened. More generally, any type of labeling, as well as rules for propagating, inheriting, changing, or otherwise manipulating such labels, may be used by the coloring system **310** as contemplated herein. A suitable coloring system is described in greater detail below with reference to FIG. 4.

(82) The key management system **312** may support management of keys for the endpoint **302** in order to selectively permit or prevent access to content on the endpoint **302** on a file-specific basis, a process-specific basis, an application-specific basis, a user-specific basis, or any other suitable basis in order to prevent data leakage, and in order to support more fine-grained and immediate control over access to content on the endpoint **302** when a security compromise is detected. Thus, for example, if a particular process executing on the endpoint is compromised, or potentially compromised or otherwise under suspicion, access by that process may be blocked (e.g., with access to keys revoked) in order to prevent, e.g., data leakage or other malicious activity. A suitable key management system useful in this context is described in greater detail below with reference to FIG. 5.

(83) The heartbeat system **314** may be used to provide periodic or aperiodic information from the

endpoint **302** or other system components about system health, security, status, and so forth. The heartbeat system **314** or otherwise an endpoint health system may thus in general include a health status report system for the endpoint **302**, such as through the use of a heartbeat system or the like. A heartbeat may be encrypted or plaintext, or some combination of these, and may be communicated unidirectionally (e.g., from the endpoint **302** to the threat management facility **308**) or bidirectionally (e.g., between the endpoint **302** and the server **306**, or any other pair of system components) on any useful schedule. A suitable heartbeat system that can be used as part of the endpoint health system is described in greater detail below with reference to FIG. 6.

(84) In general, these various monitoring and management systems may cooperate to provide improved threat detection and response. For example, the coloring system **310** may be used to evaluate when a particular process is potentially opening inappropriate files, and a potential threat may be confirmed based on an interrupted heartbeat from the heartbeat system **314**. The key management system **312** may then be deployed to revoke access by the process to certain resources (e.g., keys or file) so that no further files can be opened, deleted, or otherwise modified. More generally, the cooperation of these systems enables a wide variety of reactive measures that can improve detection and remediation of potential threats to an endpoint.

(85) FIG. 4 illustrates a system for behavioral tracking, coloring, and generation of indications of compromise (IOCs). In general, the system **400** may include a number of entities participating in a threat management process such as any of the entities and threat management processes described herein. The threat management process may for example employ techniques such as behavioral tracking, encryption, endpoint recording, reputation-based threat detection, behavioral-based threat detection, signature-based threat detection, and combinations of the foregoing, or any other suitable techniques for detecting threats to endpoints in an enterprise.

(86) In general, the system **400** may include a number of endpoints **402**, **412** and a threat management facility **404** in an enterprise **410**, such as any of the enterprises described herein. An external analysis facility **406** may analyze threat data and provide rules and the like for use by the threat management facility **404** and endpoints **402**, **412** in managing threats to the enterprise **410**. The threat management facility **404** may reside locally (e.g., a part of, embedded within, or locally coupled to the endpoint **402**), a virtual appliance (e.g., which could be run by a protected set of systems on their own network system(s)), a private cloud, a public cloud, and so forth. The analysis facility **406** may store locally-derived threat information. The analysis facility **406** may also or instead receive threat information from a third-party source **416** such as MITRE Corporation or any other public, private, educational, or other organization that gathers information on network threats and provides analysis and threat detection information for use by others. Each of these components may be configured with suitable programming to participate in the various threat detection and management techniques contemplated herein. The threat management facility **404** may monitor any stream of data from an endpoint **402** exclusively or use the full context of intelligence from the stream of all protected endpoints **402**, **412** or some combination of these.

(87) The endpoint **402** may be any of the endpoints described herein, or any other device or network asset that might join or participate in the enterprise **410** or otherwise operate on an enterprise network. This may, for example, include a server, a client such as a desktop computer or a mobile computing device (e.g., a laptop computer, a wearable device, a tablet, and the like), a cellular phone, a smart phone, or other computing device suitable for participating in the enterprise **410**.

(88) In general, the endpoint **402** may include any number of computing objects such as an object **418** labeled with a descriptor **420**. While the term object has a number of specific meanings in the art, and in particular in object-oriented programming, it will be understood that the term 'object' as used herein is intended to be significantly broader, and may include any data, process, file or combination of these including without limitation any process, application, executable, script, dynamic linked library, file, data, database, data source, data structure, function, resource locator

(e.g., uniform resource locator (URL) or other uniform resource identifier (URI)), or the like that might be manipulated by one of the computing devices described herein.

(89) An object **418** may also or instead include a remote resource, such as a resource identified in a URL. That is, while the objects **418** in FIG. 4 are depicted as residing on the endpoint **402**, an object **418** may also reside elsewhere in the system **400**, while still being labeled with a descriptor **420** and tracked by the monitor **421** of the endpoint **402**. The object **418** may be an item that is performing an action or causing an event, or the object **418** may be an item that is receiving the action or result of an event (i.e., the item in the system **400** being acted upon).

(90) Where the object **418** is data or includes data, the object **418** may be encrypted or otherwise protected, or the object **418** may be unencrypted or otherwise unprotected. The object **418** may be a process or other computing object that performs an action, which may include a single event or a collection or sequence of events taken by a process. The object **418** may also or instead include an item such as a file or lines of code that are executable to perform such actions. The object **418** may also or instead include a computing component upon which an action is taken, e.g., a system setting (e.g., a registry key or the like), a data file, a URL, or the like. The object **418** may exhibit a behavior such as an interaction with another object or component of the system **400**.

(91) In one aspect, objects **418** may be described in terms of persistence. The object **418** may, for example, be a part of a process, and remain persistent as long as that process is alive. The object **418** may instead be persistent across an endpoint **402** and remain persistent as long as an endpoint **402** is active or alive. The object **418** may instead be a global object having persistence outside of an endpoint **402**, such as a URL or a data store. In other words, the object **418** may be a persistent object with persistence outside of the endpoint.

(92) Although many if not most objects **418** will typically be benign objects forming a part of a normal, operating endpoint, an object **418** may contain software associated with an advanced persistent threat (APT) or other malware that resides partially or entirely on the endpoint **402**. The associated software may have reached the endpoint **402** in a variety of ways and may have been placed manually or automatically on the endpoint **402** by a malicious source. It will be understood that the associated software may take any number of forms and have any number of components. For example, the associated software may include an executable file that can execute independently, or the associated software may be a macro, plug-in, or the like that executes within another application. Similarly, the associated software may manifest as one or more processes or threads executing on the endpoint **402**. Further, the associated software may install from a file on the endpoint **402** (or a file remote from the endpoint **402**), and the associated software may create one or more files such as data files or the like while executing. Associated software should be understood to generally include all such files and processes except where a specific file or process is more specifically noted.

(93) A threat such as an APT may also take the form of an attack where no altered or additional software is directly added or modified on the endpoint **402**. Instead, an adversary may reuse existing software on the system **400** to perform the attacks. It is for this reason that simply scanning for associated software may be insufficient for the detection of APTs and it may be preferable to detect APTs based on the behavior of the software and associated objects **418** that are used by, for, and with that software.

(94) An object coloring system **414** may apply descriptors **420** to objects **418** on the endpoint **402**. This may be performed continuously by a background process on the endpoint **402**, or it may occur whenever an object **418** is involved in an action, such as when a process makes a call to an application programming interface (API) or takes some other action, or when a URL is used to initiate a network request, or when a read or a write is performed on data in a file. This may also or instead include a combination of these approaches as well as other approaches, such as by pre-labeling a file or application when it is moved to the endpoint **402**, or when the endpoint **402** is started up or instantiated. In general, the object coloring system **414** may add, remove, or change a

color at any location and at any moment that can be practicably instrumented on a computer system.

(95) As noted above, the term ‘object’ as used herein is intended to include a wide range of computing objects and as such, the manner in which particular objects **418** are labeled or ‘colored’ with descriptors **420** may vary significantly. Any object **418** that is performing an action may be colored at the time of and/or with a label corresponding to the action, or likewise any object **418** that is the target of the action may be colored at the time that it is used and/or with a label corresponding to a process or the like using the object **418**. Furthermore, the operating system runtime representation of the object **418** may be colored, or the persistent object outside of the operating system may be colored (as is the case for a File Handle or File Object within the operating system or the actual file as stored in a file system), such as within an encryption header or other header applied to the file, or as part of a directory attribute or any other persistent location within the file or file system. A former coloring may be ephemerally tracked while the operating system maintains the representation and the latter may persist long after any reboots of the same operating system and likewise have meaning when read or used by other endpoints **402**. For processes, each file handle may be supplemented with a pointer or other mechanism for locating a descriptor **420** for a particular object **420** that is a process. More specifically, each object **418** may be colored in any manner suitable for appending information to that object **418** so that the corresponding descriptor **420** can be retrieved and, where appropriate, updated.

(96) The coloring system **414** may apply any suitable rules for adding and changing descriptors **420** for objects **418**. For example, when a process with a certain descriptor accesses data with a different descriptor, the descriptor for the process may be updated to correspond to the data, or the descriptor for the data may be updated to correspond to the process, or some combination of these. Any action by or upon an object **418** may trigger a coloring rule so that descriptors **420** can be revised at any relevant time(s) during processing.

(97) In one aspect, colors will not explicitly indicate a compromised security state or other good/bad types of distinctions (although they may be adapted to this use). Instead, colors may record some known information or understanding about an object **418**, such as a source, a purpose, and so forth. In this context, colors will not be used to label actual or potential security compromises, but to identify inconsistencies among interacting objects **418**, and to restrict or control access and use accordingly. For example, where an endpoint uses file-system-based encryption as described herein, a process that is colored as exposed to external resources (e.g., the Internet) may be prohibited from accessing cleartext data for protected files. Colors can also be used in other contexts such as intrusion prevention, routing rules, and detection of odd or questionable behavior.

(98) In one aspect, colors may be implemented as flags associated with objects **418** that provide a short hand cache of potentially relevant information. While this information could also be obtained for an object **418** through a careful inspection of related activity logs or other data recording activities, the use of a cache of flags for coloring information makes the coloring information directly available and immediately actionable, as distinguished from post hoc forensic activities that are otherwise supported by data logging.

(99) In one aspect, colors as contemplated herein may fall into two different categories: static colors and dynamic colors. Static colors may be explicitly applied based on, e.g., a controlling application. For example, a static color may specify a status of an application or data, or an associated type of application (e.g., productivity, mail client, messaging, browser, word processing, financial, spreadsheet, etc.). In this context, a process will generally inherit static colors from a source executable, and will permit inferences for appropriate behavior and related processes. Dynamic colors may be assigned based on direct observation of executing processes, and may not be inherited or transferred among processes (although the presence of a dynamic color may be used to draw another coloring inference upon interaction with another process). Thus, the inheritance of

colors may depend in part upon the type of color that is applied, or upon explicit inheritance rules provided for a particular color.

(100) A descriptor **420** may take a variety of forms, and may in general include any information selected for relevance to threat detection. This may, for example, be a simple categorization of data or processes such as trusted or untrusted. For example, in one embodiment described herein, data and processes are labeled as either 'IN' (e.g., trusted) or 'OUT' (e.g., untrusted). The specific content of the label is unimportant, and this may be a binary flag, text string, encrypted data or other human-readable and/or machine-readable identifier, provided that the descriptor **420** can facilitate discrimination among labeled files—in this example, between trusted objects **418** and untrusted objects **418** so that, e.g., trusted data can be selectively decrypted or encrypted for use with trusted processes. Similarly, data may be labeled as corporate data or private data, with similar type-dependent processing provided. For example, private data may be encrypted with a key exclusively controlled by the data owner, while corporate data may be encrypted using a remotely managed key ring for an enterprise operated by the corporation.

(101) In another aspect, the descriptor **420** may provide a multi-tiered or hierarchical description of the object **418** including any information useful for characterizing the object **418** in a threat management context. For example, in one useful configuration the descriptor **420** may include a type or category, static threat detection attributes, and an explicit identification. The type or category for the object **418** may be any category or the like that characterizes a general nature or use of the object **418** as inferred from behavior and other characteristics. This may, for example, include categories such as 'game,' 'financial,' 'application,' 'electronic mail,' 'image,' 'video,' 'browser,' 'antivirus,' and so forth. The category may be more granular, or may include hierarchical categories such as 'application:spreadsheet,' 'application:word processing,' and so forth. Such colors may be directly inferred from a single action, a sequence of actions, or a combination of actions and other colors, including, e.g., colors of processes and files related to a particular action, or other objects **418** that provide context for a particular action or group of actions. One or more colors may also or instead be explicitly provided by a user or a process, or otherwise automatically or manually attributed to computer objects as contemplated herein.

(102) The static threat detection attributes may be any readily ascertainable characteristics of the object **418** useful in threat detection. This may, for example, include an antivirus signature, a hash, a file size, file privileges, a process user, a path or directory, declarations of permissions, an access (e.g., a resource access, or an API access), and so forth. Static threat detection attributes may also include attributes that are derived by or supplied from other sources. For example, static threat detection attributes may include a reputation for an object **418**, which may be expressed in any suitable or useful level of granularity such as with discrete categories (trusted/untrusted/unknown) or with a numerical score or other quantitative indicator. The explicit identification may, in general, be what an object **418** calls itself, e.g., a file name or process name.

(103) Some actions may transfer colors from a subject of the action to the target of the action. For example, when a process creates sub-processes, the sub-processes may inherit the colors of its parent(s). By way of another example, when a process is initially loaded from an executable, it may inherit the color(s) stored in the file system for or with the executable.

(104) In general, the descriptor **420** may be provided in any suitable format. The descriptor **420** may for example be formed as a vector of binary flags or other attributes that form the 'color' or description of an object **418**. The descriptor **420** may also, where appropriate, include scalar quantities for certain properties. For example, it may be relevant how many times a system file was accessed, how many file handles a process has open, how many times a remote resource was requested or how long a remote resource is connected, and this information may be suitably included in the descriptor **420** for use in coloring objects with the coloring system **414** and applying rules for IOC detection by the IOC monitor **421**.

(105) An indication of compromise (IOC) monitor **421** may be provided to instrument the endpoint

402 so that any observable actions by or involving various objects **418** can be detected. As with the coloring system **414**, it will be understood that the types of observable actions will vary significantly, and the manner in which the endpoint **402** is instrumented to detect such actions will depend on the particular type of object **418**. For example, for files or the like, an API for a file system may be used to detect reads, writes, and other access (e.g., open, read, write, move, copy, delete, etc.), and may be configured to report to or otherwise initiate monitoring of the action taken with the file through the file system. As another example, kernel objects may be instrumented at the corresponding object handle or in some other manner. As a further example, a kernel driver may be used for intercepting a process startup. While a wide variety of objects are contemplated herein, one of ordinary skill in the art may create suitable instrumentation for any computing object so that it may be monitored by the IOC monitor **421**.

(106) It will be noted that suitable instrumentation may be used for a variety of functions and circumstances. For example, instrumentation may usefully track requests for network access or other actions back to a particular application or process, or data payloads back to a particular file or data location. One of ordinary skill in the art can readily implement suitable traces and/or logging for any such information that might be useful in a particular IOC monitoring operation.

(107) In general, the IOC monitor **421** applies rules to determine when there is an IOC **422** suitable for reporting to a threat management facility **404**. It will be understood that an endpoint **402** may, in suitable circumstances and with appropriate information, take immediate local action to remediate a threat. However, the monitor **421** may advantageously accumulate a sequence of actions, and still more advantageously may identify inconsistencies or unexpected behavior within a group of actions with improved sensitivity by comparing descriptors **420** for various objects **418** involved in relevant actions and events. In this manner, rules may be applied based upon the descriptors **420** that better discriminate malicious activity while reducing the quantity and frequency of information that must be communicated to a remote threat management facility **404**. At the same time, all of the relevant information provided by the descriptors **420** can be sent in an IOC **422** when communicating a potential issue to the threat management facility **404**. For example, during the course of execution, a specific process (as evidenced by its observed actions) may be assigned color descriptors indicating that it is a browser process. Further, the specific process may be assigned an attribute indicating that it has exposed itself to external URLs or other external data. Subsequently, the same process may be observed to be taking an action suitable for an internal or system process, such as opening up shared memory to another process that has coloring descriptions indicating that it is a system process. When this last action is observed, an inconsistency in the various color descriptors between the subject of the action—the externally exposed browser process—and the target of the action may result in a well-defined IOC, which may be directly processed with immediate local action taken. The IOC may also or instead be reported externally as appropriate.

(108) Thus, an endpoint **402** in an enterprise **410** may be instrumented with a coloring system **414** and monitor **421** to better detect potentially malicious activity using descriptors **420** that have been selected for relevance to threat detection along with a corresponding set of rules developed for the particular descriptors **420** that are being used to label or color various objects **418**. By way of example, the object **418** may be a web browser that starts off being colored as a ‘browser’ and an ‘internet facing’ application. Based on this descriptor **420**, a range of behaviors or actions may be considered normal, such as accessing remote network resources. However, if an object **418** colored with this descriptor **420** attempted to elevate privileges for a process, or to access a registry or system files, then this inconsistency in action may trigger a rule violation and result in an IOC **422**.

(109) In general, any action or series of actions that cumulatively invoke a particular reporting or action rule may be combined into an IOC **422** and communicated to the threat management facility **404**. For example, an IOC **422** may include a malicious or strange behavior, or an indication of a malicious or strange behavior. The IOC **422** may be a normalized IOC that expresses one or more

actions in a platform independent manner. That is, the IOC **422** may express a malicious behavior or suspected malicious behavior without reference to platform-specific information such as details of an operating system (e.g., iOS, MacOS, Windows, Android, Linux, and so forth), hardware, applications, naming conventions, and so forth. Thus, a normalized IOC may be suitable for identifying a particular threat across multiple platforms, and may include platform independent processes, actions, or behaviors, or may express such process, actions, or behaviors in a platform independent manner. The normalized IOC may be generated from the IOC **422**, e.g., it may be a converted version of the IOC **422** suitable for use with multiple platforms, or it may simply be any IOC **422** that has been created in a platform independent form. Process colorization (i.e., using the coloring system **414**) as described herein may be used to create a normalized IOC.

(110) In general, a threat management facility **404** for the enterprise **410** may include an IOC collector **426** that receives the IOC **422** from the endpoint **402** and determines an appropriate action. This may include any suitable remedial action, or where one or more IOCs **422** are inconclusive, continued monitoring or increased monitoring as appropriate.

(111) The threat management facility **404** may provide a variety of threat management or monitoring tools **424**, any of which may be deployed in response to IOCs **422** collected by the IOC collector **426**. These tools **424** may include without limitation a scanning engine, whitelisting/blacklisting, reputation analysis, web filtering, an emulator, protection architecture, live protection, runtime detection, APT detection, network antivirus products, IOC detection, access logs, a heartbeat, a sandbox, or quarantine system, and so forth.

(112) The analysis facility **406** may provide a remote processing resource for analyzing malicious activities and creating rules **434** suitable for detecting IOCs **422** based on objects **420** and descriptors **420**. It is generally contemplated that suitable attributes of certain descriptors **418** and one or more rules **434** may be developed together so that objects **418** can be appropriately labeled with descriptors **420** that permit invocation of rules **434** and creation of IOCs **422** at appropriate times. The analysis facility **406** may include a variety of analysis tools **428** including, without limitation, tools for regular expression, whitelisting/blacklisting, crowd sourcing, identifiers, and so forth. The analysis tools **428** may also or instead include information and tools such as URL look-ups, genotypes, identities, file look-up, reputations, and so forth. The analysis facility **406** may also provide numerous related functions such as an interface for receiving information on new, unknown files or processes, and for testing of such code or content in a sandbox on the analysis facility **406**.

(113) The analysis facility **406** may also or instead include a compromise detector **430**, where the compromise detector **430** is configured to receive new threat information for analysis and creation of new rules and descriptors as appropriate, as well as corresponding remedial actions. The compromise detector **430** may include any tools described herein or otherwise known in the art for detecting compromises or evaluating new threats in an enterprise **410**.

(114) In general, a rule **434** may be manually created with corresponding human-readable semantics, e.g., where a process is labeled as a browser process or other category or type that can be interpreted by a human. It should, however, be appreciated that the compromise detector **430** may also be configured to automatically generate descriptors **420** and rules **434** suitable for distribution to a threat management facility **404** and an endpoint **402**. In this latter mode, the meaning of a particular descriptor **420** may not have a readily expressible human-readable meaning. Thus, it will be understood that attributes selected for relevance to threat detection may include conventional attributes, as well as attributes without conventional labels or meaning except in the context of a particular, computer-generated rule for threat detection.

(115) In general, the analysis facility **406** may be within an enterprise **410**, or the analysis facility **406** may be external to the enterprise **410** and administered, for example, by a trusted third party. Further, a third-party source **416** may provide additional threat data **438** or analyses for use by the analysis facility **406** and the threat management facility **404**. The third-party resource **416** may be a

data resource that provides threat data **438** and analyses, where the threat data **438** is any data that is useful in detecting, monitoring, or analyzing threats. For example, the threat data **438** may include a database of threats, signatures, and the like. By way of example, the third-party resource **416** may be a resource provided by The MITRE Corporation.

(116) The system **400** may include a reputation engine **440** storing a plurality of reputations **442**. The reputation engine **440** may include a reputation management system for the generation, analysis, identification, editing, storing, etc., of reputations **442**. The reputation engine **440** may include reputation-based filtering, which may be similar to the reputation filtering discussed above with reference to FIG. **1**. The reputation engine **440** may be located on the threat management facility **404** or the endpoint **402** as shown in FIG. **4**, or the reputation engine **440** may be located elsewhere in the system **400**. The reputation engine **440** may receive an IOC **422** or a stream of IOCs **422**, and may generate or utilize reputations **442** for the IOCs **422**. The reputation engine **440** may also or instead receive actions, behaviors, events, interactions, and so forth, and may generate or utilize reputations **442** for any of the foregoing. The reputation engine **440** may generate or revise a reputation **442** based on behaviors, actions, events, interactions, IOCs **422**, other reputations **442**, a history of events, data, rules, state of encryption, colors, and so forth. The reputation engine **440** may utilize a third-party resource, e.g., for the third-party resource's reputation data.

(117) The reputations **442** may include reputations for any of the objects **418** as described herein. In general, the reputations **442** may relate to the trustworthiness of the objects **418** or an attribute thereof (e.g., the source of the object **418**, a behavior of the object **418**, another object interacting with the object **418**, and so forth). The reputations **442** may include lists of known sources of malware or known suspicious objects **418**. The reputations **442** may also or instead include lists of known safe or trusted resources or objects **418**. The reputations **442** may be stored in a reputations database included on the reputation engine **440** or located elsewhere in the system **400**. The reputations **442** may be expressed in any suitable or useful level of granularity such as with discrete categories (e.g., trusted, untrusted, unknown, malicious, safe, etc.) or with a numerical score or other quantitative indicator. The reputations **442** may also be scaled.

(118) In general, in the system **400** of FIG. **4**, a malicious activity on the endpoint **402** may be detected by the IOC monitor **421**, and a corresponding IOC **422** may be transmitted to the threat management facility **404** for remedial action as appropriate. The threat management facility **404** may further communicate one or more IOCs **422** to the analysis facility **406** for additional analyses and/or resolution of inconclusive results. Other details and variations are provided below. While the use of coloring and IOCs as contemplated herein can improve threat detection and remediation in a number of ways, the system **400** can be further improved with granular control over access to endpoint data using an encryption system. A system for key-based management of processes and files on an endpoint is now discussed in greater detail.

(119) FIG. **5** illustrates a system for encryption management. Generally, the system **500** may include endpoints **502**, an administration host **504**, and a threat management facility **506**, which may include policy manager **508** and key manager **510**. The system **500** may provide for the management of users **512**, policies **514**, keys **516** (e.g., disposed on key rings **518**), and endpoints **502** (e.g., from the administration host **504**). The system **500** may utilize various storage and processing resources, which may be local, remote, virtual, disposed in a cloud, or the like.

(120) The endpoints **502** may be any of the endpoints as described herein, e.g., with reference to the other figures. The endpoints **502** may also or instead include other end user devices and other devices to be managed. The endpoints **502** may include a web browser for use by the users **512**, with supporting cryptographic functions implemented using cryptographic libraries in the web browser. The endpoints **502** may communicate with the other components of the system **500** using any suitable communication interface, which may include Secure Socket Layer (SSL) encryption, Hypertext Transfer Protocol Secure (HTTPS), and so forth for additional security.

(121) The endpoints **502** may include objects as described herein. For example, the endpoints **502** may include processes **520** and files **522**. The processes **520** may be labeled (e.g., by a coloring system using descriptors as described above) in such a manner that the process is ‘IN,’ where the process **520** is in compliance with policies **514** administered for the endpoint **502** from a remote threat management facility **506**, or the process is ‘OUT,’ where the process **520** is out of compliance with a policy (or a number of policies) in the policies **514** for an enterprise. This may provide IN processes **520A** and OUT processes **520B** as shown in FIG. 5. The files **522** may be similarly labeled by a coloring system with descriptors that identify each file **522** as IN, where the file **522** complies with the policies **514** and is accordingly encrypted using, e.g., a remotely managed key ring **518**, or the file is OUT, where the file **522** does not conform to the policies **514** and is accordingly not encrypted using the remotely managed key ring **518**. This may provide IN files **522A** and OUT files **522B** as shown in FIG. 5. One skilled in the art will recognize that other objects of the endpoint **502** or other components of the system **500** may be labeled in a similar manner where they are either IN or OUT. By coloring objects in this manner and basing key access on the corresponding color, the “IN” software objects may operate in a protected environment that objectively appears to be in compliance with the policies **514**. Other files and processes may still be used on the endpoint **502**, but they will operate in an “OUT” or unprotected environment that cannot obtain access to any of the “IN” content or functionality.

(122) In an implementation, the system **500** may include determining whether an endpoint **502** is IN or OUT or whether a component of the endpoint **502** is IN or OUT, which may be based upon a set of rules (e.g., the rules outlined herein) or policies such as the policies **514** described herein. In some aspects, if the entire endpoint **502** is OUT—that is, out of compliance with one or more policies **514**, the endpoint **502** will not have key access or access to any protected content.

Conversely, if the endpoint **502** is IN, the endpoint **502** may have access to protected content. Thus, in one aspect, the notion of IN/OUT may be applied at an endpoint level, and data protection may be a consequence of endpoint protection. Endpoint protection may also or instead be applied at a more granular level, e.g., by determining whether executables, processes **520**, files **522**, etc., on the endpoint **502** are IN or OUT, which may be based upon rules or policies **514** as described herein.

(123) The administration host **504** may include a web browser, which may include a cryptography library **524** and a web user interface (e.g., HTML, JavaScript, etc.). An administrator may utilize the web user interface to administer a key management system and perform administrative functions such as creating and distributing keys **516**, establishing security policies, creating key hierarchies and rules, and so forth. The endpoint **502** may also include a cryptographic library **524** implementing cryptographic protocols for using key material in the key ring **518** to encrypt and decrypt data as needed.

(124) The threat management facility **506** may include any of the threat management facilities or similar systems described herein. In general, the threat management facility **506** may include a policy manager **508** and key manager **510**. Alternatively, one or more of the policy manager **508** and key manager **510** may be located elsewhere on a network.

(125) The policy manager **508** may implement one or more policies **514**, and maintain, distribute, and monitor the policies for devices in an enterprise. The policies **514** may include any policies **514** relating to secure operation of endpoints **502** in an enterprise. This may, for example, include hardware configuration policies, software configuration policies, communication policies, update policies, or any other policies relating to, e.g., the configuration of an endpoint **502**, communications by an endpoint **502**, software executing on an endpoint **502** and so forth. Policies **514** may include usage criteria based on, e.g., signatures, indications of compromise, reputation, user identity, and so forth. With respect to the key management system contemplated herein, the policies **514** may include a cryptographic protocol design, key servers, user procedures, and other relevant protocols, or these cryptographic protocols may be provided elsewhere for use by the policy manager **508**. The policies **514** may also include any rules for compliance including those

mentioned above or any other suitable rules or algorithms that can be applied to determine whether objects and components are 'IN' or 'OUT' as contemplated herein.

(126) The key manager **510** may be part of the threat management facility **506**, or it may be remotely managed elsewhere, e.g., in a remote cloud resource or the like. The key manager **510** may also or instead be disposed on the administration host **504** and one or more endpoints **502** in a manner independent of the threat management facility **506**. In this manner, all cryptographic operations may be isolated from the threat management facility **506** and instead may be performed by a web browser or the like executing on the administration host **504** or an endpoint **502**. The key manager **510** may manage the keys **516**, including managing the generation, exchange, storage, use, and replacement of keys **516**. The key manager **510** may include a key ring **518**, where the keys **516** are disposed on the key ring **518** using one root key **526**. The key manager **510** may also or instead include a variety of key management and other secure processes, including without limitation, administrator registration, establishing trust to endpoints **502**, key distribution to endpoints **502**, policy deployment, endpoint status reporting, and local key backup.

(127) The users **512** may have full access to encrypted data. Alternatively, the users **512** may have limited access to encrypted data, or no access to encrypted data. Access may be limited to users **512** using endpoints **502** that are deemed 'IN' by the system, as well as to processes **520** that are IN, as further described herein.

(128) The keys **516** may include cryptographic keys in a cryptosystem, i.e., decryption keys. In one aspect, the keys **516** may be disposed on one key ring **518** using one root key **526**. In general, the keys **516** may be created and managed using, e.g., symmetric key technology, asymmetric key technology, or any other key technology or combination of key technologies suitable for securing data in an enterprise including, for example the Data Encryption Standard (DES), Triple DES, Advanced Encryption Standard (AES), elliptic curve cryptography (ECC), and so forth. The cryptosystem may also or instead include any suitable public key infrastructure or the like supporting the distribution and use of keys for encryption, digital signatures, and so forth.

(129) The key ring **518** may facilitate simplified management of the system **500**. For example, by reducing the data protection system down to a single key ring **518**, the system can eliminate or reduce the overhead for management of keys **516**. In one aspect, all of the data on a key ring **518** is protected by one root key **526**. By reducing the data protection system down to a single key ring **518** protected by one root key **526**, all privileged users **512** on uncompromised platforms can have access to all protected data. In this embodiment, data is either 'IN' (i.e., encrypted), or it is 'OUT' (i.e., not encrypted). In one aspect, the default system does not include any additional level of granularity of access control.

(130) The cryptography library **524** may be disposed on the administration host **504** as shown in FIG. 5. The cryptography library **524** may also be disposed on the endpoint **502**, e.g., in a web browser, or it may be disposed on another component of the system **500**, or any combination of these. The cryptographic library **524** may be installed by an administrator. In general, key material **530** from the key ring **518** may be stored in a cache **532** on the endpoint **502** within any suitable memory on the endpoint **502** for use in encryption and decryption as contemplated herein. As noted above, an enterprise that systematically uses coloring and indications of compromise can be improved through the use of a synchronized or integrated key management system as contemplated herein. This system may be still further improved with the addition of a heartbeat system that communicates heartbeats from an endpoint containing health and status information about the endpoint. A suitable heartbeat system is now described in greater detail.

(131) FIG. 6 illustrates a threat management system using heartbeats. In general, a system **600** may include an endpoint **602**, a gateway **604**, a threat management system **606**, and an enterprise management system **608** that manages an enterprise including the endpoint **602**, the gateway **604**, and one or more additional endpoints **610**. Each of these components may be configured with suitable programming to participate in the detection and remediation of an advanced persistent

threat (APT) or other malware threat as contemplated herein. Although the term “gateway” is used for the device between an endpoint and an external network, it will be appreciated that this device may also or instead include a switch, router, firewall, and/or other network elements, any of which may be included in the “gateway” as that term is used herein.

(132) The endpoint **602** may be any of the endpoints described herein, or any other device or network asset that might join or participate in an enterprise network. The endpoint **602** may contain a threat **612** such as an advanced persistent threat, virus, or similar malware that resides on the endpoint **602**. The threat **612** may have reached the endpoint **602** in a variety of ways, and may have been placed manually or automatically on the endpoint **602** by a malicious source. It will be understood that the threat **612** may take any number of forms and have any number of components. For example, the threat **612** may include an executable file that can execute independently, or the threat **612** may be a macro, plug-in, or the like that executes within another application. Similarly, the threat **612** may manifest as one or more processes or threads executing on the endpoint **602**. The threat **612** may install from a file on the endpoint **602** or a file remote from the endpoint **602**, and the threat **612** may create one or more other files such as data files or the like while executing. Advanced persistent threats can be particularly difficult to detect and remediate, and the systems and methods contemplated herein can advantageously provide improved sensitivity to such threats, as well as enabling improved remediation strategies. However, the systems and methods contemplated herein may also or instead be used to detect and remediate other types of malware threats. As such, in this context references to a particular type of threat (e.g., an advanced persistent threat) should be understood to generally include any type of malware or other threat to an endpoint or enterprise unless a more specific threat or threat type is explicitly provided or otherwise clear from the context.

(133) The threat **612** may be analyzed by one or more threat countermeasures on the endpoint **602** such as a whitelisting filter **614** that approves each item of code before executing on the endpoint **602** and prevents execution of non-whitelisted code. The endpoint **602** may also include an antivirus engine **616** or other malware detection software that uses any of a variety of techniques to identify malicious code by reputation or other characteristics. A runtime detection engine **618** may also monitor executing code to identify possible threats. More generally, any of a variety of threat detection techniques may be applied to the threat **612** before and during execution. In general, a threat **612** may evade these and other security measures and begin executing as a process **620** on the endpoint **602**.

(134) Network traffic **622** from the process **620** may be monitored and logged by a traffic monitor **624** on the endpoint **602**. The traffic monitor **624** may, for example, log a time and a source of each network request from the endpoint **602**. Where the endpoint **602** is within an enterprise network, the network traffic **622** may pass through the gateway **604** in transit to a data network such as the Internet. While the gateway **604** may be logically or physically positioned between the endpoint **602** and an external data network, it will be understood that other configurations are possible. For example, where the endpoint **602** is associated with an enterprise network but operating remotely, the endpoint **602** may form a VPN or other secure tunnel or the like to the gateway **604** for use of a threat management system **606**, enterprise management system **608**, and any other enterprise resources.

(135) The endpoint **602** may use a heartbeat **626** to periodically and securely communicate status to the gateway **604**. The heartbeat **626** may be created by a health monitor **628** within the endpoint **602**, and may be transmitted to a remote health monitor **630**, for example, at the gateway **604**. The health monitor **628** may monitor system health in a variety of ways, such as by checking the status of individual software items executing on the endpoint **602**, checking that antivirus and other security software is up to date (e.g., with current virus definition files and so forth) and running correctly, checking the integrity of cryptographic key stores, checking for compliance with enterprise security policies, and checking any other hardware or software components of the

endpoint **602** as necessary or helpful for health monitoring. The health monitor **628** may thus condition the issuance of a heartbeat **626** on a satisfactory status of the endpoint **602** according to any suitable criteria, enterprise policies, and other evaluation techniques. The remote health monitor **630** may also or instead be provided at the threat management facility **650**, for example as part of the threat management system **606** or the enterprise management system **608**.

(136) The heartbeat **626** may be secured in any suitable manner so that the health monitor **630** can reliably confirm the source of the heartbeat **626** and the status of the endpoint **602**. To this end, the heartbeat **626** may be cryptographically signed or secured using a private key so that the monitor **630** can authenticate the origin of the heartbeat **626** using a corresponding public key. In one aspect, the heartbeat **626** may include a combination of plaintext information and encrypted information, such as where the status information for the endpoint is provided in plaintext while a digital signature for authentication is cryptographically secured. In another aspect, all of the information in the heartbeat **626** may be encrypted.

(137) In one aspect, a key vault **632** may be provided on the endpoint to support cryptographic functions associated with a secure heartbeat. An obfuscated key vault **632** may support numerous useful functions, including without limitation, private key decryption, asymmetric signing, and validation with a chain of trust to a specific root validation certificate. A variety of suitable key management and cryptographic systems are known in the art and may be usefully employed to support the use of a secure heartbeat as contemplated herein. The system may support a secure heartbeat in numerous ways. For example, the system may ensure that signing and decryption keys can only be used in authorized ways and inside an intended Access Control mechanism. The system may use “anti-lifting” techniques to ensure that a signing key can only be used when the endpoint is healthy. The system may ensure that attacking software cannot, without first reverse-engineering the key vault **632**, extract the original key material. The system may also usefully ensure that an attacker cannot undetectably replace the public keys in a root certificate store, either directly or indirectly, such as in an attack that tries to cause the code to validate against a different set of root keys without directly replacing any keys in the root store.

(138) A robust heartbeat **626** may usefully provide defensive mechanisms against reverse engineering of obfuscated content (e.g., the private key material stored in key vault **632**, the code used to validate the correct running of the remainder of the systems as part of the heartbeat **626** code itself) and any anti-lifting protections to prevent malware from directly using the endpoint **602** (or the health monitor **628** on the endpoint **602**) to continue to send out signed heartbeat packets (e.g. stating that “all is well” with the endpoint) after security mechanisms have been impaired, disabled, or otherwise compromised in any way. Lifting in this manner by malicious code can be materially mitigated by providing statistical validation (e.g., with checksums of code) of call stacks, calling processes, and core processes. Likewise, statistical checks as well as checksum integrations into the cryptographic calculations may protect against code changes in the heartbeat **626** code itself.

(139) A variety of useful techniques may be employed to improve security of the key vault **632** and the heartbeat **626**. For example, the system may use domain shifting so that original key material is inferred based on hardware and software properties readily available to the key vault **632**, and to ensure that key material uses non-standard or varying algorithms. Software properties may, for example, include readily determined system values such as hashes of nearby code. In another aspect, the keys may be domain shifted in a manner unique to the endpoint **602** so that the manner of statistical validation of call stacks and core software is unique to the endpoint **602**. Further the key vault may be provisioned so that a public key stored in the key vault **632** is signed with a certificate (or into a certificate chain) that can be externally validated by a network appliance or other trusted third party or directly by the health monitor **628** or remote health monitor **630**.

(140) The heartbeat **626** may encode any useful status information, and may be transmitted from the endpoint **602** on any desired schedule including any periodic, aperiodic, random, deterministic,

or other schedule. Configured in this manner, the heartbeat **626** can provide secure, tamper-resistant instrumentation for status of the endpoint **602**, and in particular an indication that the endpoint **602** is online and uncompromised. A delay or disappearance of the heartbeat **626** from the endpoint **602** may indicate that the endpoint **602** has been compromised; however, this may also simply indicate that the endpoint **602** has been powered off or intentionally disconnected from the network. Thus, other criteria may be used in addition to the disappearance or interruption of the heartbeat **626** to more accurately detect malicious software. Some such techniques are described below, but it will be understood that this may include any supplemental information that might tend to make an attack on the endpoint **602** more or less likely. For example, if the heartbeat **626** is interrupted but the endpoint **602** is still sourcing network traffic, then an inference might suitably be made that the endpoint **602** is compromised.

(141) The threat management system **606** may, in general, be any of the threat management systems described herein. The enterprise management system **608** generally provides tools and interfaces for administration of the enterprise and various endpoints **610** and other resources or assets attached thereto. It will be understood that the functions of the threat management system **606** and the enterprise management system **608** may vary, and general threat management and administration functions may be distributed in a variety of ways between and among these and other components. This is generally indicated in FIG. **6** as a threat management facility **650** that includes the threat management system **606** and the enterprise management system **608**. It will be understood that either or both of these systems may be administered by third parties on behalf of the enterprise, or managed completely within the enterprise, or some combination of these, all without departing from the scope of this disclosure. It will similarly be understood that a reference herein to a threat management facility **650** is not intended to imply any particular combination of functions or components, and shall only be understood to include such functions or components as explicitly stated in a particular context, or as necessary to provide countermeasures for malware (e.g., advanced persistent threats) as contemplated herein. It also should be understood that the heartbeat may be monitored and/or managed by the threat management system **606**, the enterprise management system **608**, or another component of the threat management facility **650**.

(142) The system **600** may include a certificate authority **660** or similar trust authority or the like (shown as a “trusted third party” in the figure). In order to provide a meaningfully secure heartbeat **626**, the heartbeat **626** may be secured with reference to a trusted authority such as a certificate authority **660** that can issue cryptographic certificates allowing other entities to rely on assertions about identity (e.g., by enabling verification with a trusted third party), and to enable cryptographically secure communications. The cryptographic techniques for creating and using such certificates and relationships are well known, and are not repeated here. The certificate authority **660** may be administered by the enterprise management system **608** or some other internal resource of the enterprise, or the certificate authority **660** may be administered by a trusted third party such as any of a variety of commercially available certificate authorities or the like. Thus, the certificate authority **660**, or some other similar cloud service or the like, may operate as a security broker to register, e.g., endpoints **602**, **610**, the gateway **604**, the threat management facility **650**, and so forth, and provide cryptographic material for each of the other trusting entities to securely communicate with one another.

(143) Once registered with the certificate authority **660** in this fashion, the heartbeat may be used to establish trust between the endpoint **602** and other entities, and to validate the source of the heartbeat **626** when it is received. More generally, a heartbeat **626** secured in this manner may provide an encrypted channel between network entities such as an endpoint **602** and the gateway **604** (or a firewall or the like). The nature of the communication may provide a technique for validating the source, as well as obfuscating the contents with encryption. Thus when, for example, the endpoint **602** provides information about a good/healthy state or a bad/compromised state, the recipient may rely on this state information and act accordingly.

(144) FIG. 7 shows an architecture for endpoint protection in an enterprise network security system. In general, an endpoint may include a processing environment **702**, a file system **706** (such as a data storage system or the like), a threat monitor **720** and a key wrapper **730**.

(145) The processing environment **702** may, for example, be any environment such as an operating system or the like suitable for executing one or more processes **704**.

(146) Each process **704** may be an instance of a software application, computer program, portion of a computer program or other code executing within the processing environment **702**. A process **704** may execute, e.g., on a processor, group of processors, or other processing circuitry or platform for executing computer-executable code. A process **704** may include executable computer code, as well as an allocation of memory, file descriptors or handles for data sources and sinks, security attributes such as an owner and any associated permissions, and a context including the content of physical memory used by the process **704**. A process **704** may be or may include one or more threads. More generally, a process **704** may include any code executing on an endpoint such as any of the endpoints described herein.

(147) The file system **706** may include a data storage system or the like, e.g., where a data store including one or more files (e.g., the files **708** shown in the figure) is included as part of the data storage system. The file system **706** may be generally associated with an operating system that provides the processing environment **702**, and serves as an intermediary between processes **704** executing in the processing environment **702** and one or more files **708** accessible to the endpoint. The file system **706** may provide a directory structure or other construct to facilitate organization of the files **708**, and the file system **706** generally supports file functions such as creating, deleting, opening, closing, reading, writing, and so forth.

(148) An extension **710** may be included in the file system **706** by modifying the operating system kernel. While other programming techniques may be employed to perform the functions of an extension **710** as contemplated herein, direct modifications to or additions to the operating system permit the extension **710** to operate transparently to the processing environment **702** and the processes **704** without requiring any modifications or adaptations. The extension **710** may, for example, be implemented as a file system filter (in a MICROSOFT WINDOWS environment) or a mount point to a directory (in an APPLE IOS environment). The extension **710** to the files system as contemplated herein performs two concurrent functions. First, the extension **710** communicates with a threat monitor **720** in order to receive updates on the security status and exposure status of the processes **704** or the endpoint. Second the extension **710** communicates with a key wrapper **730** that provides key material for encrypting and decrypting data in the files **708**. Finally, the extension **710** operates to conditionally provide encryption and decryption of the files **708** for the processes **704** based on a current security or exposure state, as described in greater detail below.

(149) The threat monitor **720** may include any suitable threat monitoring, malware detection, antivirus program or the like suitable for monitoring and reporting on a security state of an endpoint or individual processes **704** executing thereon. This may include local threat monitoring using, e.g., behavioral analysis or static analysis. The threat monitor **720** may also or instead use reputation to evaluate the security state of processes **704** based on the processes **704** themselves, source files or executable code for the processes **704**, or network activity initiated by the processes **704**. For example, if a process **704** requests data from a remote URL that is known to have a bad reputation, this information may be used to infer a compromised security state of the endpoint. While a threat monitor **720** may operate locally, the threat monitor **720** may also or instead use remote resources such as a gateway carrying traffic to and from the endpoint, or a remote threat management facility that provides reputation information, malware signatures, policy information and the like for the endpoint and other devices within an enterprise such as the enterprise described above.

(150) The threat monitor **720** may also or instead monitor the health of one or more of the system, an endpoint, a process **704**, and so forth. The health monitoring may be used to provide periodic or

aperiodic information from one or more system components about system health, security, status, and so forth. Implementations may include using the health monitoring for controlling access, e.g., to files **708**, to keys **734**, to key material for encrypting and decrypting individual files **708**, and so forth.

(151) In general, the threat monitor **720** provides monitoring of a security state and an exposure state of the endpoint. The security state may, for example, be ‘compromised’, ‘secure’, or some other state or combination of states. This may be based on detections of known malware, suspicious activity, policy violations and so forth. The exposure state may be ‘exposed’ or ‘unexposed’, reflecting whether or not a particular process **704** or file **708** has been exposed to potentially unsafe content. Thus, exposure may not necessarily represent a specific threat, but the potential for exposure to unsafe content. This may be tracked in a variety of ways, such as by using the coloring system described above with reference to FIG. 5.

(152) The key wrapper **730** may contain a key ring **732** with one or more keys **734** for encrypting and decrypting files **708**. The key ring **732** may be cryptographically protected within the key wrapper **730** in order to prevent malicious access thereto, and the key wrapper **730** may communicate with the extension **710** to provide keys **734** for accessing the files **708** at appropriate times, depending, for example, on whether processes **704** are secure or exposed. In one aspect, the files **708** are stored in a non-volatile memory such as a disk drive, or in a random-access memory that provides a cache for the disk drive, and the key wrapper **730** may be stored in a separate physical memory such as a volatile memory accessible to the operating system and the extension **710** but not to processes **704** executing in the user space of the processing environment **702**.

(153) In one aspect, every document or file on the endpoint may have a separate key. This may be, for example, a unique, symmetric key that can be used for encryption and decryption of the corresponding file. The key wrapper **730** may control access to the key material for encrypting and decrypting individual files, and may be used by the extension **710** to control access by individual processes **704** executing on the endpoint. As described herein, the extension **710** may generally control access to files **708** based on an exposure state, a security state, or other context such as the user of a calling process or the like. In the event of a severe compromise, or a detection of a compromise independent of particular processes, a key shredding procedure may be invoked to destroy the entire key wrapper **730** immediately and prevent any further access to the files **708**. In such circumstances, the keys can only be recovered by the endpoint when a remediation is confirmed. Alternatively, the files may be accessed directly and decrypted from a secure, remote resource that can access the keys **734**.

(154) FIG. 8 illustrates a system for forensic analysis for computer processes. The system **800** may include an endpoint **810** containing a data recorder **820**, a monitoring facility **830**, and any number of objects **812** and events **814**. An analysis facility **840** may be coupled in a communicating relationship with the endpoint **810** over a data network **850** such as any of the networks described above. It will be appreciated that, while illustrated as components of the endpoint **810**, certain components of the system **800** such as the data recorder **820** and the monitoring facility **830** and the analysis facility may also or instead be realized as remote services instantiated on a virtual appliance, a public or private cloud, or the like, any of which may be coupled to the endpoint **810** through the data network **850** or another communication channel (not shown). Each of the components of the system **800** may be configured with suitable programming and configuration to participate in the various forensic techniques, threat detection techniques, and security management techniques contemplated herein.

(155) The endpoint **810** may be any of the endpoints described herein, e.g., a computing device in an enterprise network, or any other device or network asset that might join or participate in an enterprise or otherwise operate on an enterprise network. This may, for example, include a server, a client device such as a desktop computer or a mobile computing device (e.g., a laptop computer or a tablet), a cellular phone, a smart phone, or other computing device suitable for participating in the

system **800** or in an enterprise.

(156) In general, the endpoint **810** may include any number of computing objects **812**, which may for example, be processes executed by one or more processors or other processing circuitry, files or data stored in memory, or any other computing objects described herein. While the term object has a number of specific meanings in the art, and in particular in object-oriented programming, it will be understood that the term ‘object’ as used herein is intended to be significantly broader, and may include any data, process, file or combination of these including without limitation any process, application, executable, script, dynamic linked library (DLL), file, data, database, data source, data structure, function, resource locator (e.g., uniform resource locator (URL) or other uniform resource identifier (URI)), or the like that might be resident on the endpoint **810** and manipulated by the endpoint **810** or another component of the system **800** or other systems described elsewhere herein. The object **812** may also or instead include a remote resource, such as a resource identified in a URL. That is, while the object **812** in the figure is depicted as residing on the endpoint **810**, an object **812** may also reside elsewhere in the system **800**, for example with a link, pointer, or reference.

(157) The object **812** may be an item that is performing an action or causing an event **814**, or the object **812** may be an item that is receiving the action or is the result of an event **814** (e.g., the object **812** may be an item in the system **800** being acted upon by an event **814** or another object **812**). In general, an event **814** as contemplated herein may be any data flow, execution flow, control flow, network flow, or other similar action or event that might causally relate objects **812** to one another. Where the object **812** is data or includes data, the object **812** may be encrypted or otherwise protected, or the object **812** may be unencrypted or otherwise unprotected. The object **812** may be a process or other computing object that performs an action, which may include a single event **814** or a collection or sequence of events **814** taken by a process. The object **812** may also or instead include an item such as a file or lines of code that are executable to perform such actions. The object **812** may also or instead include a computing component upon which an action is taken, e.g., a system setting (e.g., a registry key or the like), a data file, a URL, and so forth. The object **812** may exhibit a behavior such as an interaction with another object or a component of the system **800**.

(158) Objects **812** may be described in terms of persistence. The object **812** may, for example, be a part of a process, and remain persistent as long as that process is alive. The object **812** may instead be persistent across an endpoint **810** and remain persistent as long as an endpoint **810** is active or alive. The object **812** may instead be a global object having persistence outside of an endpoint **810**, such as a URL or a data store. In other words, the object **812** may be a persistent object with persistence outside of the endpoint **810**.

(159) Although many if not most objects **812** will typically be benign objects forming a normal part of the computing environment for an operating endpoint **810**, an object **812** may contain software associated with an advanced persistent threat (APT) or other malware that resides partially or entirely on the endpoint **810**. This associated software may have reached the endpoint **810** in a variety of ways, and may have been placed manually or automatically on the endpoint **810** by a malicious source. It will be understood that the associated software may take any number of forms and have any number of components. For example, the associated software may include an executable file that can execute independently, or the associated software may be a macro, plug-in, or the like that executes within another application. Similarly, the associated software may manifest as one or more processes or threads executing on the endpoint **810**. Further, the associated software may install from a file on the endpoint **810** (or a file remote from the endpoint **810**), and the associated software may create one or more files such as data files or the like while executing. Associated software should be understood to generally include all such files and processes except where a specific file or process is more specifically noted.

(160) An event **814** may include an action, a behavior, an interaction, and so forth. The event **814**

may be generated by or otherwise related to an object **812**. For example, the event **814** may be associated with a file and include an action such as a read, a write, an open, a move, a copy, a delete, and so forth. The event **814** may also or instead include an inter-process communication, e.g., a create, a handle, a debug, a remote injection, and so forth. The event **814** may also or instead include accessing an Internet Protocol (IP) address or URL.

(161) The data recorder **820** may monitor and record activity related to the objects **812** and events **814** occurring on the endpoint **810**. The activity of the endpoint **810** may be stored in a data log **822** or the like on the data recorder **820**, which may be stored locally on the endpoint **810** (as depicted) or remotely at a threat management resource, or some combination of these, such as where the data log **822** is periodically transmitted to a remote facility for archiving or analysis. The data recorder **820** may continuously record any activity occurring on the endpoint **810** for predetermined periods of time before overwriting previously recorded data. Thus, the data log **822** may include a continuous data feed of events **814**. When an event **814** is detected that is a beacon or trigger event (such as a file detection, a malicious traffic detection, or the like), the data log **822** may be saved and transmitted to an analysis facility **840** or the like for analysis, e.g., to determine a root cause of the beacon or trigger event. The data log **822** may be used to create an event graph or other snapshot of the activity on the endpoint **810**, e.g., for a period of time surrounding a beacon or trigger event. The beacon or trigger event may be detected locally by the monitoring facility **830**, or remotely by a remote threat management facility or the like, or some combination of these.

(162) While illustrated on the endpoint **810**, it will be understood that the data recorder **820** may also or instead be implemented at a remote location such as a threat management facility or other enterprise network security resource. The data recorder **820** may be provisioned on the same or a different device than a data store in which data is stored. The data recorder **820** may be configured to record data as efficiently as possible so as to minimize impact on the endpoint **810**.

(163) The monitoring facility **830** may work in conjunction with the data recorder **820** to instrument the endpoint **810** so that any observable events **814** by or involving various objects **812** can be monitored and recorded. It will be appreciated that various filtering rules and techniques may be used to synopsise, summarize, filter, compress or otherwise process information captured by the data recorder **820** to help ensure that relevant information is captured while maintaining practical limits on the amount of information that is gathered.

(164) A security product **832** may execute on the endpoint **810** to detect a security event on the endpoint **810**, which may act as the beacon or trigger event for the system **800**. The security product **832** may use techniques such as signature-based and behavioral-based malware detection including without limitation one or more of host intrusion prevention, malicious traffic detection, URL blocking, file-based detection, and so forth.

(165) The beacon or trigger event on the endpoint **810** may be a fully qualified (e.g., definitive) detection of a compromise or other malicious activity. In another aspect, the beacon or trigger event on the endpoint **810** may be a suspicious behavior that is suspicious but not confirmed as malicious. For example, the beacon or trigger event on the endpoint **810** may signal an unusual behavior that is known to commonly appear concurrently with the detection of malware. In an aspect, when the beacon or trigger event is a suspicious behavior, the data log **822** may be analyzed differently than when the beacon or trigger event is a confirmed malicious behavior. For example, the data log **822** may be sent to a different component of the system **800** through the network, e.g., to a different analysis facility **840**.

(166) The monitoring facility **830** may be disposed remotely from the endpoint **810** or analysis facility **840**. The monitoring facility **830** may be included on one or more of the endpoint **810** or analysis facility **840**. In an aspect, the monitoring facility **830** and the analysis facility **840** included in the same component.

(167) The analysis facility **840** may analyze the data log **822**, e.g., as part of a root cause analysis and to identify objects **812** compromised by the root cause. To this end, the analysis facility **840**

may utilize one or more rules **842** for applying to the data included in the data log **822** to determine a root cause of a beacon or trigger event such as a suspected or actual security compromise on the endpoint **810**. The analysis facility **840** may reside locally on the endpoint **810** (e.g., be a part of, embedded within, or locally coupled to the endpoint **810**). The analysis facility **840** may be an external facility, or it may reside in a virtual appliance (e.g., which could be run by a protected set of systems on their own network systems), a private cloud, a public cloud, and so forth. The analysis facility **840** may store locally-derived threat information for use in subsequent identification, remediation, or other similar activity. The analysis facility **840** may also or instead receive threat information from a third-party source such as any public, private, educational, or other organization that gathers information on network threats and provides analysis and threat detection information for use by others. This third-party information may, for example, be used to improve detection rules or other forensic analysis that might be performed on information in the data log **822**.

(168) The analysis facility **840** may create an event graph. In general, the event graph may represent information in the data log **822** in a graph where objects **812** are nodes and events **814** are edges connecting the nodes to one another based on causal or other relationships as generally contemplated herein. The event graph may be used by the analysis facility **840** or other component(s) of the system **800** as part of a root cause analysis and to identify objects **812** compromised by the root cause. The event graph may also or instead be displayed to a user of the system **800** or endpoint **810**, e.g., using an interactive user interface or the like.

(169) The system **800** may advantageously use the data log **822** to configure and initialize an analysis in a sandboxed or otherwise isolated environment where the execution of the recorded activity related to a detected security event is allowed to run. That is, rather than uploading a complete image of an endpoint **810** using conventional techniques, the data log **822** may include only a series of events/processes related to the detected event that may be uploaded for execution/analysis. The analysis may thus include executing this series of events/processes in the same order to determine a threat level for the endpoint **810**.

(170) The data log **822** may include data from a single endpoint **810**, or from a number of endpoints **810**, for example where one endpoint **810** accesses a service or a file on another endpoint. This advantageously facilitates tracking or detection of potentially malicious activity that spans multiple devices, particularly where the behavior on a single endpoint does not appear malicious. Thus, the monitoring facility **830** may monitor activity from an endpoint **810** exclusively, or use the full context of activity from all protected endpoints **810**, or some combination of these. Similarly, the event graph generated from the data log **822** may include activity from one endpoint **810** exclusively, or use the full context of activity from all protected endpoints **810**, or some combination of these. Data logs **822** and event graphs may also or instead be stored for future analyses, e.g., for comparing to future data logs and event graphs.

(171) FIG. **9** is a flowchart of a method for forensic analysis for computer processes. The method **900** may be implemented by any of the systems described above or otherwise herein. The method **900** may be used as part of a root cause analysis, e.g., for determining a root cause of malware on an endpoint, and for identifying computing objects affected by malware, e.g., computing objects causally related to the root cause.

(172) As shown in step **902**, the method **900** may include monitoring events on a device, such as a first endpoint. The events may be any as described herein, e.g., events associated with computing objects on the endpoint. The computing objects may, for example include a data file, a process, an application, a registry entry, a network address, a peripheral device, or any of the other computing objects described herein. For example, in an aspect, the computing objects may include one or more network addresses specified at any suitable level of abstraction or according to any suitable protocol such as a uniform resource locator (URL), an Internet Protocol (IP) address, and a domain name, and may include any or a portion of associated path information or the like that might be

associated therewith. The computing objects may also or instead include a peripheral device such as a universal serial bus (USB) memory, a camera, a printer, a memory card, a removable bulk storage device, a keyboard, a printer, a scanner, a cellular phone, or any other input or output device that might usefully be connected to an endpoint, a server, a mobile device, and so forth. Events may include information or messages from a threat management facility, firewall, network device, and so on, for example, that may be resident on or in communication with an endpoint. For example, a threat management facility may identify a potential or actual threat, and this may be treated as an event.

(173) In an aspect, monitoring events on a first endpoint may include instrumenting a first endpoint to monitor a number of causal relationships among a number of computing objects. For example, a monitoring facility or other monitoring component (e.g., a component disposed on the first endpoint or otherwise in communication with the first endpoint), may be configured to detect computing objects and to monitor events on the first endpoint that associate the computing objects in a number of causal relationships. Thus, a processor and a memory disposed on the endpoint may be configured to monitor events on the endpoint. A remote server may also or instead be configured to monitor events on the endpoint, for example, to create a data log as contemplated herein.

(174) Implementations may also or instead include monitoring events on multiple endpoints, e.g., endpoints included in an enterprise network or the like. Thus, in an aspect, the one or more computing objects include at least one or more computing object(s) on a device other than the first endpoint, such as a second endpoint in the enterprise network. The device may also or instead include a server configured to provide remote resources to other endpoints, network devices, firewalls, gateways, routers, wireless access points, mobile devices, and so forth.

(175) The causal relationships monitored by the system may include dependencies that form a link or an association between computing objects or events. Useful causal relationships may include a data flow, e.g., linking computing objects based on the flow of data from one computing object to another computing object. The causal relationships may also or instead include a control flow. For example, a first computer program may generate a first event that triggers a second computer program to trigger a second event, thereby creating a causal relationship between the first computer program and the second computer program (and possibly a causal relationship between the first event and the second event). In yet another aspect, the causal relationships may include a network flow. For example, a computing object may access a URL or other remote resource or location and receive data. In this example, there may be a causal relationship between one or more of the computing object, the URL, and the data. It will be understood that the term “causal relationship” and the like is intended to cover a wide range of relationships between computing objects that might be formed by events, and unless explicitly stated to the contrary or otherwise clear from the text, the causal relationships may include anything that can link or associate multiple computing objects (of the same type or different types), e.g., in a directional manner, directly or indirectly.

(176) As shown in step **904**, the method **900** may include recording events such as any of the events described above that occur on the endpoint. Thus, each event detected during monitoring may be recorded, e.g., by a data recorder or other component, to provide a data log including a sequence of events causally relating the number of computing objects. As described above, the data recorder may be configured to record events that occur on the endpoint, or events that occur on a plurality of endpoints. The data recorder may be locally disposed on the endpoint or otherwise in communication with the endpoint. The data recorder may also or instead be associated with a monitoring facility or an analysis facility such as any of those described above. The data recorder may record a sequence of events causally relating a number of computing objects on one or more endpoints in a data log or the like disposed in a memory.

(177) A number of events within the sequence of events may be preserved for a predetermined time window. For example, in an aspect, a data recorder or the like may record all activity on an endpoint in a rolling buffer that overwrites data that is older than the predetermined time window.

This may be true regardless of the types of computing objects associated with the sequence of events. In another aspect, the predetermined time window may have a different duration for different types of computing objects (e.g., for at least two types of computing objects). By way of example, when the computing objects include one or more network addresses, the sequence of events may be preserved for a longer predetermined time window relative to a sequence of events associated with data files, or vice-versa. Similarly, when the computing objects include one or more peripheral devices such as USB memories, the sequence of events may be preserved for longer predetermined time window relative to a sequence of events associated with applications, or vice-versa. In implementations, the predetermined time window for which the sequence of events is preserved may be based on the likelihood of a security event originating from a certain type of computing object. For example, the reputation of a computing object (e.g., an application) or a machine state may be used for determining the duration of the predetermined time window for which the sequence of events is preserved. Further, the predetermined time window for which the sequence of events is preserved may be determined by a color of a computing object or event, e.g., as described in U.S. patent application Ser. No. 14/485,759 filed on Sep. 14, 2014, which is incorporated by reference herein in its entirety. In an aspect, the time window for which the sequence of events is preserved may be variable or adjustable. For example, a user or administrator using a user interface or the like may adjust the time window for which the sequence of events is preserved, e.g., based on computing object type or otherwise. For example, one or more first event types may be recorded with a first time window and one or more second event types may be recorded with a second time window.

(178) In an aspect, the data recorder or the like may record only certain activity on an endpoint, e.g., activity associated with predetermined computing objects. The activity may be preserved for a predetermined amount of time dependent upon the specific computing object to which the activity is associated. In this manner, and by way of example, the data recorder or the like may include a record of data for one week for applications, for three months for files, for two weeks for registry entries, and so forth. It will be understood that these timeframes are provided by way of example and not of limitation.

(179) In general, data may be continuously recorded, periodically recorded, or some combination of these. Furthermore, data may be cached, stored, deleted, or transmitted to a remote processing facility in any suitable manner consistent with appropriate use of local and remote resources, and the utility or potential utility of information that is being recorded. In one aspect, data may be periodically deleted or otherwise removed from the data recorder, such as after a security event has been detected and addressed as described below. A new data log may then be created for recording subsequent events on the one or more endpoints.

(180) As shown in step **906**, the method **900** may include evaluating one or more events that occur on the endpoint. The evaluation of the one or more events may include the application of one or more security rules to determine whether the one or more events indicate or suggest a security event such as a security compromise event, a data exposure, a malware detection, or the like. Thus, the evaluation of the one or more events may lead to the detection of a security event. While illustrated as a separate step, this step **906** may be performed concurrently with or in sequence with the monitoring step **902** discussed above.

(181) The security event may be any beacon or trigger event, such as any of those discussed herein. The security event may include an event that is related to network security, computer security, data security, data leakage, data exposure, or any other actual or potential security issue. The security event may also or instead include other events of interest that are not directly related to computer/network security where, for example, they are useful for otherwise auditing or monitoring machines or characterizing device behavior. Thus, the security event may be any event general related to operation of a computer, and does not necessarily include an actual security compromise event. However, in implementations, the security event may include an actual

compromise to a network, an endpoint, or a computer system such as the detection of malware or any other threat detection. For example, the security event may be a security compromise event related to a specific threat, e.g., an event related to computer-based malware including without limitation a virus, spyware, adware, a Trojan, an intrusion, an advanced persistent threat, spam, a policy abuse, an uncontrolled access, and so forth.

(182) Detecting the security event may include detecting a security compromise by applying a static analysis to software objects on the first endpoint. For example, each software object may be individually analyzed for its compliance with a security policy or the like using signatures or other objective characteristics. It will be understood that while static analysis provides one useful form of evaluation for compliance with the security policy or the like, other techniques may also or instead be employed, e.g., a behavioral analysis, a sandbox execution, network traffic analysis, and so forth.

(183) Detecting the security event may also or instead include detecting a security compromise by applying dynamic or behavioral analysis to code executing on the first endpoint, or to specific computing objects (e.g., processes) on the endpoint. For example, events that can warrant triggering the detection of the security event may include a process that loads a particular file that is known to be malicious, or a process that accesses a known malicious IP address, and the like.

(184) In an aspect, detecting the security event may include detecting a hardware change or other state changes. Detecting the security event may also or instead include detecting a potential data leakage.

(185) As discussed herein, a security policy may be used to detect a security event. This may include, for example, whitelists or blacklists of known computing objects and events, or reputations and signatures thereof. For example, a security policy may include rules that allow computing objects and events that are provided by a known, trusted source (e.g., a trusted user, endpoint, network, company, vendor, and so forth). The rules may be more complex, for example, where originating from a trusted source is only one factor in determining whether to whitelist computing objects and events. In general, the security policy may include any suitable rules, logic, prioritizing, etc., as desired to detect a security event.

(186) Although referred to herein in terms of ‘security,’ one skilled in the art will recognize that a security policy may also or instead include other types of policies. For example, a security policy may include a corporate or network policy having a list of approved computing objects and events, where computing objects and events outside of this list may not necessarily be security risks, but are otherwise unwanted in the network. Thus, the security policy may intend to detect malware and the like, while also detecting other types of unwanted computing objects and events that do not qualify as malware.

(187) More generally, any technique or combination of techniques suitable for evaluating endpoint activity for the detection of actual or potential security compromises may be used to detect security events as contemplated herein.

(188) As shown in step **908**, if a security event is not detected, the method **900** may return to step **902** where monitoring can continue. As further shown in step **908**, if a security event is detected, a root cause analysis or the like may be performed to identify a source of the security event as further described below. That is, detecting a security event associated with one of the number of computing objects may trigger further analysis of other causally related computing objects on an endpoint (or in certain cases, remote from an endpoint) to identify a cause of the security event, as distinguished from the symptom that generated the beacon or trigger for the analysis.

(189) As shown in step **910**, the method **900** may include generating an event graph. The event graph may be generated in response to detecting the security event, e.g., using the data log from the data recorder. The event graph may be generated at the same time as or as part of creating the data log. The event graph may include the sequence of events causally relating the number of computing objects, and more specifically, the sequence of events and computer objects causally associated

with the object(s) that triggered the detected security event.

(190) As discussed herein, the event graph may be generated based on a data log of events and computer objects stored by a data recorder during operation of the endpoint. In particular the data recorder may provide a dump of logged activities, which may be causally associated into a graph for analysis, navigation, display and so forth. Any useful portion of the data log may be used. For example, the data recorder may provide event data for a window of time before, after or surrounding the detected security event. The data log may be filtered, e.g., when the data is written to the data log (for example, by aging events as described above) or when the event graph is generated, or some combination of these. A variety of filtering techniques may be usefully employed. For example, certain types of objects or events may be removed from an event graph for specific trigger events, or certain groups of events may be condensed into a single event, such as all normal activity that occurs when a user logs into an endpoint. Similarly, computing objects that are too remote, either within the event graph or timewise, may be pruned and removed, particularly if they have a known, low diagnostic significance. Thus, the event graph may be filtered and condensed in a variety of manners to obtain a useful snapshot of events optimized for root cause analytics. Filtering of the data may be dependent upon the type of security event that is detected. Filtering of the data may adjust the level of detail included in the event graph based on memory limits, user parameters, security event type, or any other object metrics or inputs. In an aspect, the data is filtered based on reputation or the like, e.g., of computing objects included therein. For example, if an application has a good reputation, the application may not include a high level of detail associated therewith in a filtered version of the data log.

(191) In one aspect, the event graph may be generated based on a data log from a number of different endpoints and thus may represent a causal chain of data from various different endpoints. This approach advantageously permits an analysis using data that spans multiple endpoints or other network devices within a single data structure or package, thus permitting identification of a root cause even when an attack employs a complex, multi-hop approach to network assets that might otherwise evade detection. Event graphs may also or instead be generated separately for different endpoints and presented to a user or analytical system as separate, discrete entities. Event graphs for endpoints may be compared with one another, e.g., as part of the root cause analysis. For example, by analyzing and comparing similar event graphs or event graphs sharing similar computing objects or events, a heuristic approach may be developed for identifying suspicious events and computing objects for one or more endpoints. Similarly, event graphs for different endpoints in the same network enterprise may be compared or combined, e.g., where two or more endpoints have been exposed to a security event or threat. For example, event graphs for similar time periods of two or more endpoints may be ascertained and analyzed.

(192) In an aspect, cross-correlating between different data logs or event graphs may be utilized in a root cause analysis. For example, if the same security event or root cause is identified on different endpoints, the endpoints may be flagged for review or remediation. This type of analysis may be used on different endpoints throughout a network.

(193) Implementations may include a number of different event graphs stored in a data store that can be used together to detect, prevent, or determine the root causes for suspicious activity or other activity of interest, e.g., a security event. As discussed herein, the event graphs may be filtered before being stored in the data store, which can remove system activity that is not of interest in such analyses. The event graphs may be searchable, e.g., for analysis of event graphs including similar computing objects or events. The event graphs may also or instead be linked to one another, e.g., event graphs including similar computing objects or events. The event graphs may be presented to a user on a user interface or the like, e.g., an interactive user interface that allows a user to see similar or related event graphs, search the event graphs, link between event graphs, and so forth.

(194) An event graph may use a conventional structure of nodes (computing objects) and events

(edges) to represent causal relationships among computing objects. This permits the use of a wide range of graph-based techniques to assist in analysis of the context leading up to a detected event. At the same time, numerous other data structures, computer representations, and visual representations of such interrelated objects and events are also known in the art, any of which may be employed as an event graph as contemplated herein, provided that enough descriptive data about the context of an endpoint is captured to facilitate the various types of analysis and response contemplated herein.

(195) As shown in step **912**, the method **900** may include, in response to detecting the security event, traversing the event graph based on the sequence of events in a reverse order from the one of the computing objects associated with the security event to one or more preceding ones of the computing objects. In general, the reverse order is a causally reverse order. For example, where a network flow, data flow or control flow has a direction from one computing object to another computing object, the reverse order will follow this flow or causal link from the receiving computing object backward toward the source computing object. However, this may also or instead include a chronological flow, such as in a complex event graph where the time of receipt for two different inputs from two different sources is relevant. In general, a review of each of the preceding computing objects may be conducted by working backward from the computing object associated with the security event, e.g., to determine a root cause of the security event. In an aspect, this may include a static analysis of each of the preceding computing objects, or a dynamic analysis of object and event interactions, or some combination of these.

(196) As shown in step **914**, the method **900** may include applying one or more rules to the computing objects preceding the security event. For example, the method **900** may include applying a cause identification rule to the preceding ones of the computing objects and the causal relationships while traversing the event graph in order to identify one of the computing objects as a cause of the security event. In general, the root cause analysis may attempt to identify a pattern in the event graph using cause identification rules to identify one of the computing objects (or a group of the computing objects and events) as a root cause of the security event.

(197) The cause identification rule may associate the cause with one or more common malware entry points. For example, common entry points include a word processing application, an electronic mail application, a spreadsheet application, a browser, or a universal serial bus (USB) drive that is attached to an endpoint, and any of these computing objects, when encountered in an event graph, may be identified as a root cause. For example, when traversing the event graph in a reverse order from the security event, if the analysis identifies an electronic mail application that opened an attachment, this may be identified as the root cause because this is often a source of compromised security on an endpoint. Similarly, when traversing the event graph in a reverse order from the security event, if the analysis identifies a USB drive, or an unsecure or unencrypted USB drive, from which a file was opened, this may be identified as a likely cause of the security event. In one aspect, multiple candidate root causes may be identified using the cause identification rules, and a final selection may be based on other contextual information such as reputation, source, etc.

(198) Security events may also or instead be caused by a certain combination of events or combinations of events and computing objects. For example, in an aspect, the cause identification rule may associate the cause of the security event with a combination that includes a first process invoking a second process and providing data to the second process. As used herein, invoking may be interpreted broadly, e.g., where any two processes share data through an intermediate file, or narrowly, e.g., where a first process specifically spawns the second process as a child process. More generally, invoking a process as used herein is intended to broadly include any causal relationship between to processes including, e.g., spawning a process, hijacking a process (e.g., seizing control of an existing process through thread injection, process hollowing, and the like), remotely launching a process over a network, instrumenting a service in the operating system, and the like. A cause identification rule may specify a particular type of invocation relationship between

two processes, or multiple types of invocation, or any relationship between two processes. Providing data from a first process to a second process may include creating a file for use by the second process. For example, the cause of a security event may include a first process that writes a file and then takes control of a second process that reads data from the file so that the first process and the second process share data through the file.

(199) Another example of a security event may include a known non-malicious application (e.g., a commonplace word processing application) launching a command line script, which may be identified as a cause of a security event. The activity underlying events that are generated may not necessarily be malicious, but they could lead to security events or other events of interest to be further analyzed. Thus, in one aspect, a cause identification rule may flag this behavior as a root cause of a security event, or as an event that is otherwise of diagnostic interest.

(200) As shown in step **916**, the method **900** may include traversing the event graph forward from an identified or presumed cause of the security event to identify one or more other ones of the computing objects affected by the cause. In this manner, an analysis of each of the computing objects in the event graph may be conducted by working forward from the root cause to other causally dependent computing objects that might be compromised or otherwise affected by the root cause. This may include labeling or otherwise identifying the potentially compromised objects, e.g., for remediation or further analysis. A pruning step may also be employed, e.g., where any computing objects that are not causally dependent on the root cause in some way are removed from the event graph.

(201) As shown in step **918**, the method **900** may include remediating one or more computing objects affected by the cause of the security event. Remediation may include deleting computing objects from the endpoint, or otherwise remediating the endpoint(s) using computer security techniques such as any described herein. In another aspect, the identification of the root cause may be used to create new detection rules capable of detecting a security event at a point in time (or causation) closer to the root cause within the event graph. Other remediation steps may include forwarding the event graph, or a filtered and pruned event graph, to a remote facility for analysis. This data may usefully provide a map for identifying sources of malware, or for ensuring thorough remediation by identifying all of the potentially compromised computing objects that should be examined after the compromise has been addressed.

(202) FIG. **10** illustrates a graphical depiction of a portion of an example event graph **1000**. The event graph **1000** may include a sequence of computing objects causally related by a number of events, and which provide a description of computing activity on one or more endpoints. The event graph **1000** may be generated, for example, when a security event **1002** is detected on an endpoint, a gateway, or a communications server, and may be based on an identification of a malicious action in real-time or in near real-time, or on a data log or similar records obtained by an event data recorder during operation of the endpoint, gateway, or communications server. The event graph **1000** may be used to determine a root cause **1004** of the security event **1002** as generally described above. The event graph **1000** may also or instead be continuously generated to serve as, or be a part of, the data log obtained by the data recorder. In any case, an event graph **1000**, or a portion of an event graph **1000** in a window before or around the time of a security event, may be obtained and analyzed after a security event **1002** occurs determine its root cause **1004**. The event graph **1000** depicted in the figure is provided by way of example only, and it will be understood that many other forms and contents for event graphs **1000** are also or instead possible. It also will be understood that the figure illustrates a graphical depiction of an event graph **1000**, which may be stored in a database or other suitable data structure.

(203) By way of example, the event graph **1000** depicted in the figure begins with a computing object that is a USB device **1012**, which may be connected to an endpoint. Where the USB device **1012** includes a directory or file system, the USB device **1012** may be mounted or accessed by a file system on an endpoint to read contents. The USB device **1012** may be detected **1013** and

contents of the USB device **1012** may be opened **1014**, e.g., by a user of the endpoint. The USB device **1012** may include one or more files and application, e.g., a first file **1016**, a second file **1018**, and a first application **1020**. The first file **1016** may be associated with a first event **1022** and the second file may be associated with a second event **1024**. The first application **1020** may access one or more files on the endpoint, e.g., the third file **1026** shown in the figure. The first application **1020** may also or instead perform one or more actions **1028**, such as accessing a URL **1030**. Accessing the URL **1030** may download or run a second application **1032** on the endpoint, which in turn accesses one or more files (e.g., the fourth file **1034** shown in the figure) or is associated with other events (e.g., the third event **1036** shown in the figure).

(204) In the example provided by the event graph **1000** depicted in the figure, the detected security event **1002** may include the action **1028** associated with the first application **1020**, e.g., accessing the URL **1030**. By way of example, the URL **1030** may be a known malicious URL or a URL or network address otherwise associated with malware. The URL **1030** may also or instead include a blacklisted network address that although not associated with malware may be prohibited by a security policy of the endpoint or enterprise network in which the endpoint is a participant. The URL **1030** may have a determined reputation or an unknown reputation. Thus, accessing the URL **1030** can be detected through known computing security techniques.

(205) In response to detecting the security event **1002**, the event graph **1000** may be traversed in a reverse order from a computing object associated with the security event **1002** based on the sequence of events included in the event graph **1000**. For example, traversing backward from the action **1028** leads to at least the first application **1020** and the USB device **1012**. As part of a root cause analysis, one or more cause identification rules may be applied to one or more of the preceding computing objects having a causal relationship with the detected security event **1002**, or to each computing object having a causal relationship to another computing object in the sequence of events preceding the detected security event **1002**. For example, other computing objects and events may be tangentially associated with causally related computing objects when traversing the event graph **1000** in a reverse order-such as the first file **1016**, the second file **1018**, the third file **1026**, the first event **1022**, and the second event **1024** depicted in the figure. In an aspect, the one or more cause identification rules are applied to computing objects preceding the detected security event **1002** until a cause of the security event **1002** is identified.

(206) In the example shown in the figure, the USB device **1012** may be identified as the root cause **1004** of the security event **1002**. In other words, the USB device **1012** was the source of the application (the first application **1020**) that initiated the security event **1002** (the action **1028** of accessing the potentially malicious or otherwise unwanted URL **1030**).

(207) The event graph **1000** may similarly be traversed going forward from one or more of the root cause **1004** or the security event **1002** to identify one or more other computing objects affected by the root cause **1004** or the security event **1002**. For example, the first file **1016** and the second **1018** potentially may be corrupted because the USB device **1012** included malicious content. Similarly, any related actions performed after the security event **1002** such as any performed by the second application **1032** may be corrupted. Further testing or remediation techniques may be applied to any of the computing objects affected by the root cause **1004** or the security event **1002**.

(208) The event graph **1000** may include one or more computing objects or events that are not located on a path between the security event **1002** and the root cause **1004**. These computing objects or events may be filtered or 'pruned' from the event graph **1000** when performing a root cause analysis or an analysis to identify other computing objects affected by the root cause **1004** or the security event **1002**. For example, computing objects or events that may be pruned from the event graph **1000** may include a USB drive and the USB device being detected **1013**.

(209) It will be appreciated that the event graph **1000** depicted in FIG. 10 is an abstracted, simplified version of actual nodes and events on an endpoint for demonstration. Numerous other nodes and edges will be present in a working computing environment. For example, when a USB

device is coupled to an endpoint, the new hardware will first be detected, and then the endpoint may search for suitable drivers and, where appropriate, present a user inquiry of how the new hardware should be handled. A user may then apply a file system to view contents of the USB device and select a file to open or execute as desired, or an autorun.exe or similar file may be present on the USB device that begins to execute automatically when the USB device is inserted. All of these operations may require multiple operating system calls, file system accesses, hardware abstraction layer interaction, and so forth, all of which may be discretely represented within the event graph **1000**, or abstracted up to a single event or object as appropriate. Thus, it will be appreciated that the event graph **1000** depicted in the drawing is intended to serve as an illustrative example only, and not to express or imply a particular level of abstraction that is necessary or useful for root cause identification as contemplated herein.

(210) The event graph **1000** may be created or analyzed using rules that define one or more relationships between events and computing objects. The C Language Integrated Production System (CLIPS) is a public domain software tool intended for building expert systems, and may be suitably adapted for analysis of a graph such as the event graph **1000** to identify patterns and otherwise apply rules for analysis thereof. While other tools and programming environments may also or instead be employed, CLIPS can support a forward and reverse chaining inference engine suitable for a large amount of input data with a relatively small set of inference rules. Using CLIPS, a feed of new data can trigger a new inference, which may be suitable for dynamic solutions to root cause investigations.

(211) An event graph such as the event graph **1000** shown in the figure may include any number of nodes and edges, where computing objects are represented by nodes and events are represented by edges that mark the causal or otherwise directional relationships between computing objects such as data flows, control flows, network flows and so forth. While processes or files are common forms of nodes that might appear in such a graph, any other computing object such as an IP address, a registry key, a domain name, a uniform resource locator, a command line input or other object may also or instead be designated to be a node in an event graph as contemplated herein. Similarly, while an edge may be formed by an IP connection, a communication, a file read, a file write, a process invocation (parent, child, etc.), a process path, a thread injection, a registry write, a domain name service query, a uniform resource locator access and so forth other edges may be designated. As described above, when a security event is detected, the source of the security event may serve as a starting point within the event graph **1000**, which may then be traversed backward to identify a root cause using any number of suitable cause identification rules. The event graph **1000** may then usefully be traversed forward from that root cause to identify other computing objects that are potentially tainted by the root cause so that a more complete remediation can be performed. The event graph **1000** may include events associated with one or more different endpoints, gateways, or communications servers.

(212) FIG. **11** shows an architecture for instrumenting an endpoint. In general, an endpoint **1100** may include a memory having a user mode **1102** providing general memory for use by user applications and processes, and a kernel mode **1104** providing protected memory exclusively for use by an operating system. Instrumentation for monitoring and remediation may, for example, be hooked into existing operating code, e.g., by injection of dynamic linked libraries into operating system code early in the boot process, preferably before the kernel libraries (e.g., kernel32.dll, for a contemporary Windows operating system) are loaded. These dynamic linked libraries, labeled hmpalert.dll in FIG. **11**, may provide an interface for programmatic access, or hooks, to code for accessing the operating system cryptography services (such as crypt32.dll, a library that implements various certificate and cryptographic functions in the Windows cryptographic application programming interface) so that execution and input/output for these kernel functions can be monitored during operation of the endpoint.

(213) FIG. **12** shows a process for accessing encrypted information on an endpoint. In general,

secrets **1202** such as cookies, tokens, credentials, authentication material, encryption keys, and so forth, may be associated with a session, a user, a group, an application, or other entity or object, and may be stored in an encrypted form on an endpoint. For example, a web browser application may store cookies so that a user can interact with a secure, remote resource without manually re-authenticating. These secrets **1202** may be encrypted using a key such as a cryptographic key for the Advanced Encryption Standard (“AES”), a symmetric block cipher (including three distinct block ciphers, AES-128, AES-192 and AES-256, with different key lengths) that uses the same key for encryption and decryption.

(214) While this layer of encryption provides some security, the underlying secrets **1202** may still be exposed to an adversary who can obtain or control use of the AES key. In order to further protect against such exploits, the AES key may itself be encrypted. For example, the Windows operating system provides a cryptographic application programming interface as a built-in component, the Data Protection Application Programming Interface (DPAPI), that uses a main cryptographic key or master key **1204** derived from a user's password. The DPAPI is stateless and does not store any persistent data itself. Instead, the DPAPI simply receives plaintext and returns ciphertext or vice versa, while relying on the operating system to protect the master key and control underlying cryptographic functions through the programming interface, e.g., through a local security authentication server known as lsass.exe in the Windows operating system.

(215) In this environment, access to a stored secret, such as a cookie or session token, involves several steps. First, a request for the stored secret is received. Then a DPAPI blob—an opaque binary structure containing encrypted data—with the encrypted AES key **1206** is retrieved and transmitted to the DPAPI (or any similar operating system cryptographic resource) for decryption with the user's master key **1204**. The decrypted AES key **1208** may then be used to decrypt the cookie for use in authentication, e.g., in order to access resources without further authentication. This multi-layered technique permits seamless access by a user without generally exposing secrets or encryption keys. However, these layers of protection do not inherently prevent programmatic requests for use of the AES key, thus opening the door to so-called pass-the-cookie attacks in which cookies are maliciously obtained by presenting improper requests to the data protection application programming interface. For example, Mimikatz is an open-source application that allows users to view authentication credentials on a Windows machine, and may be used to unprotect and steal protected cookies or login credentials stored by a web browser. ChromePass is a password recovery utility that allows users to similarly view in plaintext any usernames and passwords stored by the Chrome web browser. While these tools typically require the user to be logged in to the compute instance storing the secret, or otherwise require a user to provide authentication credentials to the operating system in order to obtain the master key used by the operating system's cryptographic infrastructure, they illustrate existing tools that are available for requesting access to secrets that include a layer of operating system cryptographic protection.

(216) As a significant advantage, the following techniques can prevent an attack aimed at improperly decrypting a local secret before the attack can present a request to the operating system's decryption resources or otherwise gain access to the key that cryptographically secures the secret. It will be understood that while much of the description herein focuses on interfaces and other resources of the Windows operating system, the principles described herein may also or instead be applied to other operating systems and that references to Windows-specific computing objects should be understood to be general in nature. For example, unless otherwise indicated, an AES key should be understood to refer to any symmetric key used to encrypt local secrets, a DPAPI Master Key should be understood to refer to any key used by an operating system to encrypt/decrypt a data blob, the Data Protection Application Programming Interface should be understood to refer to any programmatic interface for accessing cryptographic tools or libraries of an operating system, and so forth. Furthermore, while symmetric keys are commonly used for the local encryption/decryption operations described herein, asymmetric keys may also or instead be

used.

(217) FIG. 13 illustrates malware attempting to access a decrypted key. A web browser or similar application can securely store a cookie associated for example with an authenticated user or session for a connection with a remote, secure resource using the techniques above and may initiate re-use of the token by requesting decryption of a key through a cryptographic resource of the operating system. One technique for improperly accessing the key is to simply request access to a cookie store for the web browser, or more specifically to request decryption of the key used to encrypt the session token (which implicitly accesses the DPAPI or similar programming interface of the operating system). By hooking the functional interface to these operating system functions, e.g., the crypt32.dll linked library of the Windows operating system, to detect requests for decryption, a new request can be assessed for possible malicious activity.

(218) In one aspect, parameters of the request or other context for the request may be evaluated for malicious activity. For example, while the browser application that uses the cookie store may properly access the cookie store to recover session tokens and other information, other applications typically cannot. Thus, a process (or application) requesting the access may be examined to determine if the process is a proper user of the cookie store (or other secure cache or repository). If access to the cookies is requested by a process other than the web browser (or some other authorized user such as a malware scanner, backup utility, or privacy cleaner), this diagnostic pair (the requesting application and the data requested) may be used to evaluate for the presence of malicious activity. If the requesting application is not the owner of the data requested or an authorized user of the data requested, then the request may be denied. While the requesting application may be identified directly based on the process name, other information may also or instead be used to identify the application, such as the code signing certificate on the requesting application, the source path for the requesting process, or an application that is explicitly associated with the requesting process. More generally, by identifying a secret that should be protected and the application(s) that are permitted to access the secret, a corresponding detection rule can be created around context observable when the instrumented programming interface (e.g., for a decryption by the operating system) is accessed.

(219) Other remedial action may also or instead be undertaken. For example, the requesting process may be scanned for malware and/or terminated. The endpoint executing the process may be scanned and/or quarantined by preventing new network connections or terminating existing network connections. Additionally or alternatively, a beacon or similar alert may be generated based on the event and transmitted to a threat management facility for automated action, or to an administrator for manual review and intervention. Additionally or alternatively, a root cause analysis may be performed to locate earlier events associated with the malicious activity, and/or to identify other processes or computing objects that might be affected.

(220) FIG. 14 illustrates malware attempting to access a decrypted key. In a system that checks the name of a requesting process or application, e.g., as described above, a secret such as an encrypted web session token may nonetheless be vulnerable to an attack where the code of a process that would be permitted to access the secret is maliciously modified with shellcode or the like injected into the process that hijacks flow control for the process. While various malware tools exist to detect shellcode creation/deployment, such as tools to prevent access to operating system API's from dynamic memory, a shellcode exploit directed at decryption of the AES key may provide additional clues to the presence of malware that may be detected and used to prevent theft of local secrets. For example, when the call to the operating system decryption resource is made, the specific calling origin in memory, or the calling code from the process may be checked, or a signature may be checked for the executable, or behavior of the process may be examined (e.g., to determine whether the request is associated with access by a user to the corresponding remote resource), and so forth.

(221) FIG. 15 illustrates malware attempting to access a decrypted key. One exemplary class of

exploit uses dynamic linked library preloading to substitute malicious libraries for legitimate ones. Where the operating system interface (e.g., ntdll.dll, for Windows) or cryptographic library is hooked (e.g., via hmpalert.dll), the calling library may be checked before access is granted to the data protection application programming interface. For example, the name of the DLL, the path for the DLL, a signature for the DLL (either a cryptographic signature associated with the DLL, or a hash or other signature useful for identifying software), or any other contextual data or the like, may be used to evaluate the calling library. If the module name or other DLL information for the calling process does not match an authorized user of the AES key, the request may be blocked and remedial action such as any of the remedial action described herein may be initiated.

(222) FIG. 16 shows a method for detecting access to an encrypted secret stored on a compute instance. In general, a data protection application programming interface or the like may be instrumented to detect when an operating system's cryptographic resources are being used to decrypt a stored secret such as a session cookie so that a security evaluation can be performed to determine whether the requesting process is an authorized user of the secret.

(223) As shown in step 1602, the method 1600 may begin with instrumenting the appropriate application programming interface. This may, for example, include instrumenting a data protection application programming interface for an operating system on an endpoint to detect access to a decryption service used by the operating system to encrypt and decrypt data blobs with a master key derived from user credentials for the endpoint. More generally, this may include instrumenting an application programming interface on a compute instance to detect access to a decryption service used by an operating system of the compute instance, or to otherwise detect access to operating system cryptographic tools used to protect secrets on the compute instance. It will be appreciated that the programming hooks for detecting calls to an operating system API are preferably installed as early in the startup of a process as possible in order to reduce opportunities for malicious intervention as each process launches.

(224) In one aspect, the application programming interface is a data protection application programming interface for the operating system, or some other interface for accessing decryption resources of the operating system. The decryption services themselves may execute in the kernel space, the user space, or some combination of these. For an operating system such as Windows, the decryption service may encrypt and decrypt using a master key derived from user credentials for the compute instance.

(225) As shown in step 1604, the method 1600 may include detecting a call from a first process executing on the compute instance to the application programming interface to unprotect a key used to cryptographically secure a secret on the compute instance. For example, this may include detecting a call from a process executing on the endpoint to the data protection application programming interface to unprotect a symmetric key used to cryptographically secure a web browser session cookie stored by a web browser application on the endpoint. More generally, the secret may include logon credentials stored in an application cache and encrypted with the key, web browser session cookie, a cryptographic key, or some other cookie, token, credential, or other protected item of information.

(226) As shown in step 1606, the method 1600 may include evaluating the calling process, for example, by comparing first process information for the first process to second process information for one or more other processes associated with the secret. For a web browser storing a web browser session cookie, this may include comparing first process information for the process to second process information for the web browser application that stored the web browser session cookie. The process information may include any process information useful for identifying the process or a source thereof. For example, the process information may include one or more of a process name, a process identifier, an application name, one or more registry keys for the process, and a path such as a path that stores executable code for the process or a path that is used by the process for other storage and retrieval during execution. In one aspect, the process information may

include a hash used to identify the process, or a digital signature or the like for verifying an identity or source of the process. In another aspect, the process information may include a DLL, a module, a calling function, a calling location, or any other information that may be used to identify the calling process or application, or to otherwise identify or verify the source of the call.

(227) In general, the process may be any process able to call the instrumented programming interface. The one or more other processes (to which the process is compared) may include any process that might usefully and legitimately store or access a secret on the compute instance. For example, the one or more other processes may include a process associated with a web browser application, a process for a privacy cleaner, a process for a backup utility, or any process associated with an authorized application that created or stored the secret.

(228) As shown in step **1608**, the method may include determining whether the calling process is an authorized user of the secret and, where appropriate, taking responsive action. For example, in response to determining that the first process is an authorized user of the secret, this may include permitting the first process to decrypt the key for use in accessing the secret and returning to step **1604** where monitoring can continue for additional calls to the application programming interface. In response to determining that the first process is not an authorized user of the secret, this may include preventing the first process from decrypting the key for use in accessing the secret and proceeding to step **1610** for possible remediation(s). In the web browser example above, this may include, in response to determining that the process is not associated with the web browser application, preventing the process from accessing the web browser session cookie with the key and initiating a remediation of the endpoint.

(229) As shown in step **1610**, the method **1600** may include, when the first process is not an authorized user, initiating a remediation. This may include any of the remediations described herein, or any other remediation or combination of remediations suitable for responding to a suspected unauthorized access to a secret stored on a compute instance. For example, the remediation may include terminating the calling process. The remediation may also or instead include quarantining the compute instance. In another aspect, the remediation may include performing a malware scan, either of the process specifically, or a directory where the code for the process is located, or the entire compute instance. In another aspect, the remediation may include generating a beacon identifying the calling process, which may be transmitted to a threat management facility for storage and response, and/or to an administrator for manual intervention, and/or to a user of the compute instance, e.g., as a display and request for intervention through a user interface of the compute instance. In another aspect, the remediation may include performing a root cause analysis to identify a source of the calling process, or any relevant preceding events. The root cause analysis may be useful in a variety of ways, including identifying a source of the malicious action, identifying other affected computing objects on the compute instance, and creating improved tools for earlier detection based on preceding events. Upon initiation or completion of any such remediation, the method **1600** may return to step **1604** where the application programming interface may be monitored for additional calls.

(230) According to the foregoing, there is described herein a compute instance comprising a memory in a user space of an operating system, the memory storing a first key encrypted with a master key derived from user credentials for the compute instance and a secret encrypted with the first key; an application programming interface configured to provide programmatic access to cryptographic tools of the operating system based on the master key; and a security function hooked to the application programming interface, the security function configured to detect a request to decrypt the first key with the application programming interface, and to determine whether a process requesting decryption of the first key is an authorized user of the secret. The secret may be a web browser session cookie and the security function may be configured to detect whether the process is associated with a web browser that stored the web browser session cookie. More generally, the secret may include any of the secrets described herein.

(231) Embodiments disclosed herein may include computer program products comprising computer-executable code or computer-usable code that, when executing on one or more computing devices, performs any and/or all of the steps thereof. The code may be stored in a non-transitory fashion in a computer memory, which may be a memory from which the program executes (such as random-access memory associated with a processor), or a storage device such as a disk drive, flash memory or any other optical, electromagnetic, magnetic, infrared, or other device or combination of devices. In another aspect, any of the systems and methods described above may be embodied in any suitable transmission or propagation medium carrying computer-executable code and/or any inputs or outputs from same.

(232) The elements described and depicted herein, including in flow charts and block diagrams throughout the figures, imply logical boundaries between the elements. However, according to software or hardware engineering practices, the depicted elements and the functions thereof may be implemented on machines through computer executable media having a processor capable of executing program instructions stored thereon as a monolithic software structure, as standalone software modules, or as modules that employ external routines, code, services, and so forth, or any combination of these, and all such implementations may be within the scope of the present disclosure. Examples of such machines may include, but may not be limited to, personal digital assistants, laptops, personal computers, mobile phones, other handheld computing devices, medical equipment, wired or wireless communication devices, transducers, chips, calculators, satellites, tablet PCs, electronic books, gadgets, electronic devices, devices having artificial intelligence, computing devices, networking equipment, servers, routers, and the like. Furthermore, the elements depicted in the flow chart and block diagrams or any other logical component may be implemented on a machine capable of executing program instructions. Thus, while the foregoing drawings and descriptions set forth functional aspects of the disclosed systems, no particular arrangement of software for implementing these functional aspects should be inferred from these descriptions unless explicitly stated or otherwise clear from the context. Similarly, it may be appreciated that the various steps identified and described above may be varied, and that the order of steps may be adapted to particular applications of the techniques disclosed herein. All such variations and modifications are intended to fall within the scope of this disclosure. As such, the depiction and/or description of an order for various steps should not be understood to require a particular order of execution for those steps, unless required by a particular application, or explicitly stated or otherwise clear from the context. Absent an explicit indication to the contrary, the disclosed steps may be modified, supplemented, omitted, and/or re-ordered without departing from the scope of this disclosure.

(233) The method steps of the implementations described herein are intended to include any suitable method of causing such method steps to be performed, consistent with the patentability of the following claims, unless a different meaning is expressly provided or otherwise clear from the context. So, for example performing the step of X includes any suitable method for causing another party such as a remote user, a remote processing resource (e.g., a server or cloud computer) or a machine to perform the step of X. Similarly, performing steps X, Y and Z may include any method of directing or controlling any combination of such other individuals or resources to perform steps X, Y and Z to obtain the benefit of such steps. Thus, method steps of the implementations described herein are intended to include any suitable method of causing one or more other parties or entities to perform the steps, consistent with the patentability of the following claims, unless a different meaning is expressly provided or otherwise clear from the context. Such parties or entities need not be under the direction or control of any other party or entity and need not be located within a particular jurisdiction.

(234) It will be appreciated that the methods and systems described above are set forth by way of example and not of limitation. Numerous variations, additions, omissions, and other modifications will be apparent to one of ordinary skill in the art. In addition, the order or presentation of method

steps in the description and drawings above is not intended to require this order of performing the recited steps unless a particular order is expressly required or otherwise clear from the context. Thus, while particular embodiments have been shown and described, it will be apparent to those skilled in the art that various changes and modifications in form and details may be made therein without departing from the spirit and scope of this disclosure and are intended to form a part of the invention as defined by the following claims, which are to be interpreted in the broadest sense allowable by law.

Claims

1. A compute instance comprising: a memory in a user space of an operating system, the memory storing: a key encrypted with a master key derived from user credentials for the compute instance and a secret encrypted with the key, and a cookie store for a web browser, the cookie store including logon credentials encrypted by the key; and a security function executing on the compute instance, wherein: the security function includes a hook to a decryption service of the operating system, the decryption service is used to access the secret by encrypting the key and decrypting the key with a master key for the operating system, the master key is derived from user credentials for the compute instance, the security function is configured to detect, with the hook, a request to access the logon credentials by decrypting the key with the decryption service, and the security function is configured to determine whether a process requesting access to the logon credentials is associated with the web browser.
2. The compute instance of claim 1, the security function further configured to, in response to determining that the process requesting access to the logon credentials is associated with the web browser, permit the process to access the logon credentials by decrypting the key with the decryption service.
3. The compute instance of claim 1, the security function further configured to, in response to determining that the process requesting access to the logon credentials is not associated with the web browser, initiate a remedial action on the compute instance.
4. The compute instance of claim 3, wherein the remedial action includes one or more of scanning the process for malware, terminating the process, and terminating a network connection for the compute instance.
5. The compute instance of claim 3, the security function further configured to identify the process requesting access to the logon credentials based on one or more of a process name for the process, a code signing certificate for an associated application, and a source path for the process.
6. A computer program product comprising computer executable code embodied in a non-transitory computer readable medium that, when executing on one or more computing devices, causes the one or more computing devices to perform the steps of: providing a hook to detect use of an interface to a decryption service of an operating system for a compute instance; detecting, with the hook, a request from a process to use the decryption service of the operating system to decrypt a key for a cookie store of a web browser on the compute instance; assessing the request for malicious activity by determining whether the process making the request is associated with the web browser; in response to determining that the process is associated with the web browser, permitting the process to decrypt the key for the cookie store; and in response to determining that the process is not associated with the web browser, initiating a remedial action.
7. The computer program product of claim 6, wherein the remedial action includes preventing the process from decrypting the key for the cookie store.
8. The computer program product of claim 6, wherein the cookie store contains one or more logon credentials.
9. The computer program product of claim 6, wherein the cookie store contains one or more web browser session cookies.

10. The computer program product of claim 6, wherein the cookie store contains a cryptographic key.
 11. The computer program product of claim 6, wherein the cookie store contains one or more tokens.
 12. The computer program product of claim 6, wherein the cookie store contains one or more items of protected information.
 13. The computer program product of claim 6, wherein the cookie store contains one or more credentials.
 14. The computer program product of claim 6, further comprising code that causes the one or more computing devices to perform the step of identifying the process making the request based on a process name for the process.
 15. The computer program product of claim 6, further comprising code that causes the one or more computing devices to perform the step of identifying the process based on a code signing certificate for an associated application.
 16. The computer program product of claim 6, further comprising code that causes the one or more computing devices to perform the step of identifying the process based on a source path for the process.
 17. The computer program product of claim 6, wherein the remedial action includes scanning the process for malware.
 18. The computer program product of claim 6, wherein the remedial action includes terminating the process.
 19. The computer program product of claim 6, wherein the remedial action includes terminating one or more network connections for the compute instance.
 20. The computer program product of claim 6, wherein the remedial action includes transmitting an alert to a threat management facility associated with the compute instance.
-