

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388844
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Toscano; Orazio

Detecting a TCP session hijacking attack

Abstract

A computing system (500) is configured to detect a Transmission Control Protocol (TCP) session hijacking attack on a TCP session. In particular, the computing system monitors for nonsequential packets of the TCP session. The nonsequential packets each comprise a sequence number (220) that is different from a next expected sequence number of the TCP session. The computing system (500) calculates, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers (220) of the nonsequential packets received in the time interval differ from the next expected sequence number. The computing system determines a variation baseline representing a trend of the variation metrics over time, and detects the TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold.

Inventors:	Toscano; Orazio (Genoa, IT)
Applicant:	Telefonaktiebolaget LM Ericsson (publ) (Stockholm, SE)
Family ID:	74125189
Assignee:	Telefonaktiebolaget LM Ericsson (publ) (Stockholm, SE)
Appl. No.:	18/039892
Filed (or PCT Filed):	December 16, 2020
PCT No.:	PCT/EP2020/086564
PCT Pub. No.:	WO2022/128088
PCT Pub. Date:	June 23, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240015170 A1	Jan. 11, 2024

Publication Classification

Int. Cl.: H04L9/40 (20220101)

U.S. Cl.:

CPC H04L63/1416 (20130101); H04L63/1466 (20130101);

Field of Classification Search

CPC: H04L (63/1416); H04L (63/1466)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6792546	12/2003	Shanklin et al.	N/A	N/A
7936687	12/2010	Parker	709/224	H04L 12/56
8239942	12/2011	Shanklin et al.	N/A	N/A
9942787	12/2017	Tillotson	N/A	H04W 24/08
2003/0051026	12/2002	Carter et al.	N/A	N/A
2005/0210533	12/2004	Copeland	726/23	H04L 63/1416
2009/0019537	12/2008	Stavrou	726/13	H04L 63/1416
2013/0166960	12/2012	Das	714/E11.159	H04L 43/50
2015/0029863	12/2014	Lai	370/237	H04L 47/27
2018/0270197	12/2017	Subramanian	N/A	H04L 63/1416
2018/0278632	12/2017	Fu	N/A	H04L 63/1416
2020/0092218	12/2019	Flomen	N/A	H04L 65/80
2023/0353221	12/2022	Djordjevic	N/A	H04B 17/318

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
104125121	12/2013	CN	N/A
106953833	12/2016	CN	N/A
109474587	12/2018	CN	N/A
WO-2008106084	12/2007	WO	H04L 63/1408

OTHER PUBLICATIONS

Barry, B. et al., "A Cross-protocol approach to detect TCP Hijacking attacks", 2007 IEEE International Conference on Signal Processing and Communications (ICSPC 2007), Nov. 24-27, 2007, Dubai, United Arab Emirates, pp. 57-60, IEEE. cited by applicant

Lopez, D. et al., "Framework for Interface to Network Security Functions", Internet Engineering

Task Force (IETF), RFC 8329, Feb. 2018, pp. 1-25, Internet Engineering Task Force (IETF). cited by applicant

Statista, “Amount of monetary damage caused by reported cyber crime to the IC3 from 2001 to 2019 (in million U.S. dollars)”, FBI; IC3: US Department of Justice, 2020, pp. 1, Statista. cited by applicant

Hares, S. et al., “Interface to Network Security Functions (I2NSF) : Problem Statement and Use Cases”, Internet Engineering Task Force (IETF), RFC 8192, Jul. 2017, pp. 1-29, Internet Engineering Task Force (IETF). cited by applicant

Du, W. “Chapter 13: Attacks on the TCP Protocol”, Internet Security: A Hands-on Approach, Second Edition, 2019, pp. 29-52, Seed. cited by applicant

Clement, J., “IC3: total damage caused by reported cyber crime 2001-2019”, Statista, Mar. 27, 2020, pp. 1-2, Statista. cited by applicant

Primary Examiner: Keehn; Richard G

Attorney, Agent or Firm: Sonoda & Kobayashi Intellectual Property Law

Background/Summary

TECHNICAL FIELD

(1) This application generally relates to the field of computer network security, and more particularly relates to detecting a Transmission Control Protocol (TCP) session hijacking attack between communication endpoints in a communication network.

BACKGROUND

(2) TCP is a session-based protocol that builds on Internet Protocol (IP) to exchange data between endpoints on a network. To establish a TCP session, the endpoints engage in a three-way handshake procedure. Once the TCP session has been established, the TCP session can be used by the endpoints to send data to each other using TCP packets.

(3) A TCP session hijacking attack traditionally involves manufacturing a TCP packet in an attempt to get an endpoint of the session to accept the TCP packet as having been genuinely sent by the other endpoint of the session. This is often called “spoofing” a TCP packet. If a spoofed packet is accepted as genuine, a malicious third party (sometimes referred to as a “man-in-the-middle”) may be able to disrupt operations on the network. For example, the spoofed packet may include commands to reboot, shut down, or delete data, among other things. Detecting that a TCP session hijacking attack is occurring is often important, or even critical, to avoiding disruption.

SUMMARY

(4) Embodiments of the present disclosure generally relate to detecting a TCP session hijacking attack.

(5) An example embodiment includes a method of detecting a TCP session hijacking attack on a TCP session. The method is implemented by a computing system. The method comprises monitoring for nonsequential packets of the TCP session. The nonsequential packets each comprise a sequence number that is different from a next expected sequence number of the TCP session. The method further comprises calculating, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers of the nonsequential packets received in the time interval differ from the next expected sequence number. The method further comprises determining a variation baseline representing a trend of the variation metrics over time. The method further comprises detecting the TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than

a threshold.

(6) In some embodiments, calculating the variation metric for each of the plurality of time intervals comprises calculating a mean difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval, or a modal difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval. In some such embodiments, the method further comprises calculating both the mean difference and the modal difference for each of the time intervals. In some such embodiments, the method further comprises comparing the mean difference and the modal difference of the given time interval to the mean difference and the modal difference of at least one time interval previous to the given time interval.

(7) In some embodiments, the plurality of time intervals further comprises one or more time intervals subsequent to the given time interval. Moreover, detecting the TCP session hijacking attack on the TCP session is further based on the variation metric in the subsequent time intervals being different from the variation baseline by less than the threshold. In some such embodiments, detecting the TCP session hijacking attack is further based on the variation metric being different from the variation baseline by less than the threshold for more than a configurable number of the time intervals.

(8) In some embodiments, the method further comprises commencing the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, more than a threshold number of the nonsequential packets was received. In other embodiments, the method further comprises commencing the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, the nonsequential packets received for the TCP session are greater in number than for one or more other TCP sessions by a threshold amount.

(9) In some embodiments, the method further comprises refraining from calculating the variation metric for one or more time intervals in which sequence number variation is different from the variation baseline by more than the threshold. In some such embodiments, the refraining from calculating the variation metric for the one or more time intervals in which the sequence number variation is different from the variation baseline by more than the threshold is responsive to the sequence number variation being different from the variation baseline for at least a configurable number of consecutive time intervals.

(10) Other embodiments include a computing system configured to monitor for nonsequential packets of a TCP session. The nonsequential packets each comprise a sequence number that is different from a next expected sequence number of the TCP session. The computing system is further configured to calculate, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers of the nonsequential packets received in the time interval differ from the next expected sequence number. The computing system is further configured to determine a variation baseline representing a trend of the variation metrics over time. The computing system is further configured to detect a TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold.

(11) In some embodiments, the computing system is further configured to perform any of the methods discussed above.

(12) In particular, in some embodiments, the computing system comprises interface circuitry configured to receive the plurality of packets of the TCP session. The computing system further comprises processing circuitry whereby the computing system is configured as noted above. The processing circuitry is communicatively coupled to the interface circuitry.

(13) In some embodiments, to calculate the variation metric for each of the plurality of time intervals, the processing circuitry is configured to calculate a mean difference between the next

expected sequence number and the sequence numbers of the nonsequential packets received in the time interval, or a modal difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval. In some such embodiments, the processing circuitry is further configured to calculate both the mean difference and the modal difference for each of the time intervals. In some such embodiments, the processing circuitry is further configured to compare the mean difference and the modal difference of the given time interval to the mean difference and the modal difference of at least one time interval previous to the given time interval.

(14) In some embodiments, the plurality of time intervals further comprises one or more time intervals subsequent to the given time interval, and the processing circuitry is further configured to detect the TCP session hijacking attack on the TCP session based on the variation metric in the subsequent time intervals being different from the variation baseline by less than the threshold. In some such embodiments, the processing circuitry is further configured to detect the TCP session hijacking attack based on the variation metric being different from the variation baseline by less than the threshold for more than a configurable number of the time intervals.

(15) In some embodiments, the processing circuitry is further configured to commence the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, more than a threshold number of the nonsequential packets was received. In other embodiments, the processing circuitry is further configured to commence the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, the nonsequential packets received for the TCP session are greater in number than for one or more other TCP sessions by a threshold amount.

(16) In some embodiments, the processing circuitry is further configured to refrain from calculating the variation metric for one or more time intervals in which sequence number variation is different from the variation baseline by more than the threshold. In some such embodiments, the processing circuitry is further configured to refrain from calculating the variation metric for the one or more time intervals in which the sequence number variation is different from the variation baseline by more than the threshold in response to the sequence number variation being different from the variation baseline for at least a configurable number of consecutive time intervals.

(17) Other embodiments include a computer program, comprising instructions which, when executed on processing circuitry of a computing system, cause the processing circuitry to carry out any of the methods discussed above.

(18) Yet other embodiments include a carrier containing such a computer program. The carrier is one of an electronic signal, optical signal, radio signal, or computer readable storage medium.

(19) Any of the embodiments discussed above may additionally or alternatively include one or more features as discussed in greater detail below.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Aspects of the present disclosure are illustrated by way of example and are not limited by the accompanying figures with like references indicating like elements. In general, the use of a reference numeral should be regarded as referring to the depicted subject matter according to one or more embodiments, whereas discussion of a specific instance of an illustrated element will append a letter designation thereto (e.g., discussion of an endpoint **110**, generally, as opposed to discussion of particular instances of endpoints **110a**, **110b**).

(2) FIG. 1 is a schematic block diagram illustrating an example network, according to one or more embodiments of the present disclosure.

- (3) FIG. 2 is a schematic block diagram illustrating an example packet, according to one or more embodiments of the present disclosure.
- (4) FIG. 3 is a schematic block diagram illustrating different categories of packets of an established TCP session, according to one or more embodiments of the present disclosure.
- (5) FIG. 4 is a flow diagram illustrating an example approach to attack detection, according to one or more embodiments of the present disclosure.
- (6) FIG. 5 is a schematic block diagram illustrating an example computing system, according to one or more embodiments of the present disclosure.
- (7) FIG. 6 is a graph that illustrates an example of sequence number variation metrics as observed and predicted over time, according to one or more embodiments of the present disclosure.
- (8) FIG. 7 is a flow diagram illustrating an example method of detecting a TCP session hijacking attack on a TCP session, according to one or more embodiments of the present disclosure.
- (9) FIG. 8 is a flow diagram illustrating an example distribution of TCP session analysis and hijacking attack detection tasks as implemented by one or more computing systems according to one or more embodiments of the present disclosure.
- (10) FIG. 9 is a schematic block diagram illustrating an example computing system, according to one or more embodiments of the present disclosure.

DETAILED DESCRIPTION

- (11) Embodiments of the present disclosure generally relate to detecting a TCP session hijacking attack. According to various embodiments, detection of a TCP session hijacking attack is based on an evaluation of the sequence numbers of packets received in order to detect an aberration that appears not to be naturally occurring. For example, packets that are frequently out-of-sequence to the same extent can serve as an indicator that an attack is occurring in some embodiments. Particular embodiments include evaluating sequence number behavior over time (e.g., over a plurality of time intervals) to look for deviances from baseline behavior. One or more of the embodiments described herein also leverage techniques for performing such evaluation in a computationally efficient manner, which can enable the embodiment to be scaled to handle relatively higher amounts of packet traffic and/or higher numbers of concurrent TCP sessions.
- (12) FIG. 1 is a schematic block diagram illustrating an example network **100**. The network **100** comprises endpoints **110a**, **110b**. The endpoints **110a**, **110b** are computing systems that exchange packets with each other using one or more established TCP sessions. In each TCP session, one of the endpoints **110a**, **110b** sends a stream of one or more packets to the other. In some embodiments, the network **100** further comprises an intermediary **130** via which one or more of the packets between the endpoints **110a**, **110b** are exchanged.
- (13) In some embodiments, the network **100** further comprises an attacker **120**. The attacker **120** in this example is a computing system that sends one or more spoofed packets to either or both of the endpoints **110a**, **110b** as part of a TCP session hijacking attack. Each spoofed packet may be sent an endpoint **110a**, **110b** directly or via the intermediary **130** in an attempt to inject the spoofed packet into the established TCP session.
- (14) Typical examples of the computing systems illustrated in FIG. 1 include a server computer, a server cluster, a personal computer, a laptop computer, a desktop computer, a workstation, a smartphone, a tablet computer, a wearable computer, a smart appliance, network attached storage, a storage area network, a router, a gateway, and/or a switch.
- (15) The network **100** further includes one or more signaling mediums (e.g., wires, cables, buses, radio carriers) that interconnect the elements of FIG. 1 and are suitable for carrying the above-discussed packets. Examples of the network **100** include (but are not limited to) one or more of: the Internet (or a portion thereof); one or more local area networks; one or more wireless networks; one or more cellular networks; one or more Ethernet networks; and/or one or more wired networks. Such a network **100** may comprise any number of other computing devices such as routers, gateways, switches, hubs, firewalls, and the like that support the exchange of such communication

signals (not shown).

(16) Although each of the elements illustrated in the example of FIG. 1 has been described as a computing system, it should be noted that one or more of the elements illustrated in FIG. 1 may, according to other embodiments, be a software entity executing on a computing system. Indeed, in at least some such alternative embodiments, one or more of the elements illustrated in FIG. 1 may execute on the same computing system, e.g., as 3GPP virtualized network functions, virtual machines, or other software entities.

(17) FIG. 2 is a schematic block diagram illustrating an example packet **200**, according to one or more embodiments of the present disclosure. The packet **200** comprises a flow identifier **210** and a sequence number **220**. The sequence number **220** is a value that identifies the position of the packet **200** within a stream of packets from one of the endpoints **110a**, **110b** to the other that are part of the TCP session. The flow identifier **210** is a value (or combination of values) that identifies a TCP session established between the endpoints **110a**, **110b**.

(18) The packet **200** may, for example, be a traditional TCP/IP packet in which an IP packet encapsulates a TCP packet. Accordingly, the packet **200** may comprise both an IP header and a TCP header (not shown in FIG. 2). For example, the packet **200** conforms to IPv6 standards. In at least some such embodiments, the flow identifier **210** may be a flow label in the IP header of the packet **200**. Alternatively, the packet **200** may conform to IPv4 standards. In at least some such embodiments, the flow identifier **210** may be a collection of header values found in the IP and TCP headers. For example, the flow identifier **210** may be a 4-tuple that comprises source and destination IP addresses comprised in the IP header, and source and destination ports comprised in the TCP header.

(19) The packet **200** may have genuinely been sent by one of the endpoints **110a**, **110b**, or it may be a spoofed packet sent by the attacker **120**. To generate a spoofed packet **200**, the attacker **120** attempts to set the sequence number **220** and flow identifier **210** to values that one of the endpoints **110a**, **110b** will accept and treat as genuine. In general, an endpoint **110** will not accept a packet **200** as genuine if the flow identifier **210** does not match an established session and the sequence number **220** is outside of a limited range of acceptable values. While an attacker **120** may not definitively know what flow identifier and sequence number values will be accepted by an endpoint **110**, the attacker **120** may attempt to guess or predict such values.

(20) FIG. 3 is a schematic block diagram that illustrates different categories of packets **200** having a flow identifier **210** of an established TCP session that may be acceptable (e.g., by an endpoint **110** or intermediary **130**) based on their sequence number **220**. In this regard, each block in FIG. 3 is associated with a corresponding sequence number. A packet **200** having a sequence number **220** that is the same as the next expected packet may raise little to no alarm, according to embodiments. However, a nonsequential packet (i.e., a packet **200** that includes a sequence number **220** that is different from the next expected sequence number) may or may not be accepted.

(21) TCP generally accommodates the reception of packets out-of-sequence, and there are legitimate reasons why a packet **200** that is genuinely part of a TCP session may be received with a sequence number **220** that differs from that of the next expected packet. For example, network congestion may cause packets **200** traversing different paths in the network **100** to experience different amounts of delay, such that a nonsequential packet **200** arrives before the next expected packet. As another example, the next expected packet may have been dropped by the network **100** at some point in transit due to network conditions.

(22) Notwithstanding, while such network conditions may explain a relatively small difference between the sequence number of a genuine received packet, such conditions generally cannot account for a sequence number difference that is quite large. Moreover, although a packet **200** is not necessarily a spoofed packet or other indicator of attack simply because the packet **200** is out-of-sequence to a large degree, it may nonetheless be more advantageous for the receiver to simply drop the packet than to buffer the packet (and occupy memory) indefinitely until the large number

of intervening packets in the stream arrive. In this regard, a computing system may maintain a buffer for retaining acceptable packets **200** having sequence numbers that fall within a particular range. These buffered packets may be removed from the buffer after they have been processed as a genuine part of the TCP stream. That said, the computing system may refrain from storing packets **200** having sequence numbers that fall outside of this range in the buffer. Indeed, such packets **200** may, e.g., be silently discarded or otherwise treated as undesirable.

(23) Because a packet **200** comprising a sequence number that is different from a next expected sequence number of the TCP session may or may not be genuine, a computing system receiving such a packet **200** may evaluate one or more additional factors to determine the manner in which such an out-of-sequence packet will be treated. In particular, embodiments of the present disclosure may evaluate whether the sequence number of the packet **200** appears to be the product of typical network conditions, or whether the sequence number appears to be a product of packet spoofing.

(24) Embodiments of the present disclosure take into account that, while normal network operations may result in TCP received sequence numbers jumps that are acceptable (e.g., if the jump is within the size of a TCP stream's receiving buffer), the N-minute integrated mean value and modal value of such jumps for genuine packets under such conditions tends not to be constant. In contrast, during a TCP session hijacking attack, the mean and/or modal value of sequence number jumps tend to be more constant than usual.

(25) One reason these sequence number jumps tend to be more constant during an attack is because a malicious attacker **120** is attempting to guess at an acceptable sequence number range that typically changes in a generally constant way. For example, the TCP stream may comprise media being streamed between endpoints **110** at a constant bitrate. Under such circumstances, the expected sequence number of that TCP stream may generally change at a correspondingly constant rate. An attacker **120** using tools to guess or predict a sequence number that is acceptable for that stream often also shifts its guesses over time at a constant rate. Thus, the attacker **120** is building into its prediction some factor that it believes represents how much the stream is genuinely changing over time, and spoofed packets produced by such predictions may tend to produce sequence number jumps that are also fairly constant. That is, attackers **120** may tend to produce packets **200** that are out-of-sequence to the same extent repeatedly. Particular embodiments of the present disclosure detect a TCP session hijacking attack by identifying this kind of sequence number behavior.

(26) FIG. 4 is a flow diagram that illustrates a staged approach **400** to attack detection in accordance with particular embodiments of the present disclosure. The approach **400** illustrated in FIG. 4 includes a packet validation stage **410** and a predictive analytics stage **420**. In some embodiments, the approach **400** further comprises a cross-correlation stage **430**.

(27) In the packet validation stage **410**, one or more checks, validations, and/or adaptations (decryptions) may be performed on received packets. For example, packets may be decrypted checksummed, and/or checked for missing or invalid values. Particular problems with a packet may be appropriately addressed, e.g., by silently discarding a packet, triggering a retransmission request, and/or triggering an alarm or other notification.

(28) The predictive analytics stage **420** may leverage regression techniques (e.g., an Autoregressive Integrated Moving Average (ARIMA)) to forecast trends that are useful for anomaly detection based on information obtained during the packet validation stage **410**. For example, the behavior of the time-varying sequence number may be stationarized and reduction in sequence numbers deviation identified in any time interval may be treated as potential or actual evidence of an attack. In this regard, the predictive analytics stage may include a learning phase **422** in which baseline behavior is determined over one or more time intervals, followed by a detection phase **424** in which deviation from the baseline may be treated as an actual or potential attack.

(29) In the cross-correlation stage **430**, information learned from the predictive analytics stage **420** with respect to a given TCP session may be cross-correlated with that of one or more other

sessions. In this regard, sequence number behavior observed in a given TCP stream that is inconsistent with that of other TCP streams may be treated as an actual or potential attack.

(30) The approach **400** may include both localized and centralized aspects in which different aspects of the approach are implemented. The example computing system **500** illustrated in FIG. 5 employs such an approach. According to the example of FIG. 5, the computing system **500** executes a plurality of virtual functions **510a-d** (e.g., one or more virtual machines and/or one or more 3GPP virtualized network functions), each of which is an endpoint **110** or intermediary **130** of one or more TCP sessions. Each of the virtual functions **510a-d** may perform one or more of the localized functions. The computing system **500** also executes a security controller **520**, which may perform one or more of the centralized functions.

(31) A virtual function **510** may perform one or more packet validation **410** functions. According to one such example, a virtual function **510** may use a dedicated counter that increments each time a nonsequential packet is received. Periodically, this jump counter may then be forwarded to the security controller **520** for use by one or more predictive analytics **420** functions. For example, the jump counters received over a plurality of periodic intervals may be used to predict future likely values of the jump counter, and a jump counter of a subsequent interval deviating from this prediction by more than a threshold amount may be treated as potential or actual evidence of an attack. More specifically, one or more jump counters may be received during the learning phase **422** of the predictive analytics **420** stage to generate the prediction for future likely values, and during the detection phase, deviations may be treated as actual or potential attacks while the security controller continues to refine its predictions each time a new jump counter value is received.

(32) This process is not limited to simply operating on the jump counter. Indeed, a similar process may be used to evaluate the extent to which the sequence numbers of nonsequential packets vary in a given interval of time based on a baseline calculated from the behavior observed in previous intervals. In this regard, sequence number jumps that tend to be less random than usual (i.e., the sequence numbers received during a given interval of time lack sufficient variation as compared to a trend established by previous intervals), the security controller **520** may treat such behavior as potential or actual evidence of an attack. Similarly, sequence number jumps that tend to be more random than usual (i.e., the sequence numbers received during a given interval of time vary to a much greater extent as compared to a trend established by previous intervals), the security controller **520** may additionally or alternatively treat such behavior as potential or actual evidence of an attack.

(33) It should be noted that, although not illustrated in FIG. 5, according to some embodiments of the computing system **500**, the virtual functions **510a-d** and the security controller **520** execute on one or more computing devices comprised in computing system **500**. Thus, embodiments of the present disclosure may be applied to distributed computing architectures in which the computing system **500** comprises distinct computing devices that are interconnected by any appropriate network (e.g., Ethernet, radio, optical) using any appropriate network topology (e.g., full mesh, star, daisy-chain).

(34) FIG. 6 is a graph that illustrates an example of sequence number variation metrics A and A', each of which is calculated based on the sequence numbers of nonsequential packets received over time. Metric A' in particular may, in some embodiments, be a second-order metric of metric A, an example of which will be explained below.

(35) The graph reflects the use of machine learning techniques according to embodiments in which regression analysis is performed to make predictions about future outcomes. Such an approach may be particularly powerful when the behavior being analyzed is generally regular and repetitive. Accordingly, the graph depicts data observed throughout a learning phase in which, e.g., such regularities or other behaviors may serve as a basis for making predictions about subsequent behavior in a later detection phase. The learning phase occurs over n periods of time (i.e., from T0

to T_n). Behavior observed during the learning phase has been used to calculate a trend that predicts expectations during a detection phase in the intervals of the detection phase subsequent to T_n .

(36) In this example, at the end of any period T_n or later, an observed metric A that is higher than the trend by a threshold $A(\text{max})$ is treated as evidence of an attack. For example, the value $A(T_n)$ may be a value that represents the number of non-sequential packets received during the period ending at T_n . It may be expected that the number of sequence number jumps in any interval of time will not be excessive, as genuine packets should generally arrive in sequence. Therefore, at time T_n , the number of sequence number jumps received during the corresponding period may be expected to be less than the threshold $A(\text{max})$. In this particular example, the threshold $A(\text{max})$ is a dynamically changing value having the same slope as the metric A value prediction trend. In some embodiments, the threshold may additionally or alternatively be dynamically configurable. Because the number of non-sequential packets received for the interval ending at time T_n in this example is less than $A(\text{max})$, the value of $A(T_n)$ may be deemed in conformance with expectations. Had the value of $A(T_n)$ been above $A(\text{max})$, the computing system **500** may enter a warning state, generate alerts, notify an administrator or take other action to protect against an actual or potential TCP session hijacking attack.

(37) FIG. **6** also shows the calculation of metric A' over time. Although this example illustrates metric A' being calculated over the same time intervals as metric A , the calculation of these metrics may be activated and/or deactivated at different times according to other embodiments. Indeed, in some particular embodiments, detection that metric A is non-conformant with expectations may trigger activation of metric A' calculations. Correspondingly, detection that metric A is within expectations may trigger deactivation of metric A' calculations. In either case, observation that behavior is non-conformant or conformant for more than a threshold number of intervals may be required to trigger activation or deactivation, respectively, of the metric calculation, according to some embodiments.

(38) In this example, if metric A' is observed to be different from the corresponding trend by less than a threshold, the computing system **500** is configured to treat this as evidence of an attack. For example, it is generally expected that the mean and/or mode of the sequence number jumps observed over a given interval of time should not stay constant when packet traffic is genuine. Therefore, in this example, a mean and/or a mode that is less than some minimum threshold difference from the trend may be treated as evidence of non-random packet spoofing. Because metric A' observed at time T_n in this example is greater than a threshold difference from the predicted trend (and in particular, less than mean $(-\text{min } \Delta)$), the value of metric $A'(T_n)$ may be deemed in conformance with expectations. Had the value of metric $A'(T_n)$ been within mean $(+\text{min } \Delta)$ and mean $(-\text{min } \Delta)$, the computing system **500** may have generated alerts, notified an administrator, or have taken other action to protect against an actual or potential TCP session hijacking attack.

(39) It should be noted that the examples shown in FIG. **6** are just a small set of examples with respect to how analytics may be performed based on sequence numbers. In particular embodiments, sequence number variation is expected to be generally non-constant, i.e., regardless of whether mean and/or mode values are considered. Accordingly, variation of less or more than a threshold amount may be treated as evidence of an attack depending on the metric evaluated. These thresholds may also vary depending on network conditions and/or one or more other factors that might influence sequence number variation. For example, a known network congestion issue may result in expected mean and/or modal values that are more constant than usual (e.g., due to an expectation of increased packet drops). Statistical metrics other than mean and mode (e.g., standard distribution, range) may also be evaluated. In short, any of the metrics described herein may be evaluated for being above, below, or within threshold constraints as may be appropriate.

(40) Having performed one or more predictive analytics **420** functions (e.g., such as that discussed above with respect to FIG. **6**), the security controller **520** may (in some embodiments) perform one

or more cross-correlation **430** functions with respect to different TCP sessions. For example, the security controller **520** may determine whether any of the predictions with respect to any stream from any of the virtual functions **510a-d** is more than a threshold number of standard deviations from the group, and if so, treat this as evidence of an actual or potential attack. Therefore, it is possible that a TCP session may not be deemed under attack based on predictive analytics performed based solely on that session's observed metrics, but may be deemed under attack when that session's metrics are compared to those of other sessions.

(41) In some embodiments, one or more functions of the predictive analytics stage **420**, the cross-correlation stage **430**, or both are activated in response to detecting potential evidence of an attack. That is, in some embodiments, more information may be required in order to determine that an attack is actually occurring. Accordingly, responsive to detecting suspicious behavior with respect to a first metric, the computing device **500** may enter a warning state in which one or more additional checks are performed (e.g., using one or more further metrics).

(42) For example, in response to detecting that the number of non-sequential packets received in a given interval has exceeded a threshold, the calculation of one or more sequence number variation metrics may commence for one or more subsequent time intervals. For example, mean and/or mode metrics may be calculated with respect to a single TCP session, and in some embodiments be cross-correlated with other TCP sessions. Aberrant behavior in one or more of these additional variation metrics may be treated as an attack on the TCP session. In some embodiments, should such aberrant behavior fail to occur within a threshold number of time intervals, calculation of these additional analytics may be suspended.

(43) One or more of the features described above may be employed in order to provide a TCP session hijacking attack detection solution that is scalable to support a large number of established TCP sessions and/or a high amount of packet traffic. For example, in some embodiments, a more rudimentary analysis is periodically performed, and responsive to that analysis triggering the computing device **500** into a warning state, further analysis is performed periodically, e.g., until the conditions that triggered the warning state abate. In particular, some embodiments may include observing that the conditions that triggered the warning state has abated for at least a threshold number of consecutive intervals before the more advanced analytics are suspended. Indeed, the thresholds and hysteresis window for may be adjusted (e.g., programmable) to be stricter or more relaxed as appropriate for the scale of the computing system **500**, e.g., to increase or decrease the frequency in which additional statistical data is generated.

(44) Moreover, one or more cross-correlation functions may compare one or more readily calculable statistical indicators (such as mean and modal values of one or more fields in TCP packets **200**) to one or more corresponding thresholds. As these computations are generally computationally inexpensive, scaling as the number of virtual functions **510** increases may be readily accomplished. This may be further advantageous in some embodiments, as the additional scale may enhance opportunities to cross-correlate the behavior of individual TCP sessions with the behavior of other TCP sessions established with other endpoints **110** and/or passing through other intermediaries **130**.

(45) Further still, particular embodiments may separate different stages of the solution for execution at either a localized or centralized location, e.g., so that when more elaborate metrics need to be computed and analyzed, these more elaborate computations are not replicated as the number of local functions increases. Rather, such second-level computations may instead be handled by a centralized security controller **520**, which may (for example) have dedicated resources for handling these elaborations without impacting virtual function performance.

(46) Further, leveraging a security controller **520** at a centralized location and having visibility to the traffic characteristics of the other network entities may enable embodiments to advantageously cross-correlate different anomalies arising in the same timeframe and involving different virtual functions **510** of the network. Among other things, this may enable the computing system **500** to

enter the warning state by a correspondence of lower thresholds without impacting overall performance. In particular, this may enhance detection time without increasing the rate of false positives. Thus, the cross-correlations performed by the security controller may be more complex and refined than what can be discerned through localized anomaly detection.

(47) In view of the above, FIG. 7 illustrates an example method **700** of detecting a TCP session hijacking attack on a TCP session. The method **700** is implemented by a computing system **500** and comprises monitoring for nonsequential packets **200** of the TCP session (block **710**). The nonsequential packets **200** each comprise a sequence number **220** that is different from a next expected sequence number of the TCP session. The method **700** further comprises calculating, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers **220** of the nonsequential packets **200** received in the time interval differ from the next expected sequence number (block **720**). The method **700** further comprises determining a variation baseline representing a trend of the variation metrics over time (block **730**). The method **700** further comprises detecting the TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold (block **740**).

(48) FIG. 8 is a flow diagram illustrating an example distribution of TCP session analysis and hijacking attack detection tasks into localized processing **800** and centralized processing **850** components, as implemented by one or more computing systems **500**. It should be noted that although the steps illustrated in FIG. 8 are largely depicted as sequential, this is for clarity of explanation. In fact, many of the elements illustrated in FIG. 8 may be performed in parallel or in a variety of other sequences in accordance with other embodiments. Further, other embodiments include any of the localized processing **800** elements performed as part of the centralized processing **850**. Correspondingly, other embodiments include any of the centralized processing **850** elements performed as part of the centralized processing **800**.

(49) In this example, packets **200** of a TCP session are an input to one or more localized processing **800** elements. In some embodiments, such localized processing **800** may be co-located (or implemented within) a virtual function **510** of the computing system **500**, and provide data used by the centralized processing **850** performed, e.g., by a security controller **520**. In general, it may be advantageous in some embodiments for localized processing **800** to be used for tasks that are relatively computationally inexpensive, whereas more complex tasks may be advantageous to perform through centralized processing **850**. In this way, as additional virtual functions **510** and TCP sessions are added to the computing system **500**, the computational burden may scale in a relatively controlled fashion.

(50) As part of the localized processing **800** of this example, in response to detecting nonsequential packets (block **810**), a jump counter is updated (block **820**) to reflect the number of nonsequential packets received. In addition, one or more variance metric(s) are computed (e.g., mean, mode) that reflect an extent to which the sequence numbers **220** of the nonsequential packets are different from the next expected sequence number of the TCP session (block **830**). One or more delta computations are performed that compare the jump counter and/or variance metric(s) calculated in a current time interval against corresponding data obtained in one or more previous intervals (block **840**).

(51) For example, the mean and mode values of sequence number jumps for received packets may be computed every N-minutes. Once computed, the difference (i.e., delta) between each of these values as compared to those of one or more previous N-minute periods may additionally be determined.

(52) These various metrics computed locally (e.g., the jump counter, variance metrics, and/or calculated differences) may be forwarded on for use in centralized processing **850**, e.g., as performed by a security controller **520** (block **845**). Based on these forwarded metrics, centralized processing may perform one or more attack detection tasks. In particular, centralized processing

850 may be updated periodically with these metrics, and detection tasks may be performed each period.

(53) In this example, centralized processing **850** include the performance of certain data validation (block **860**), predictive analytics (block **870**), and cross-correlation (block **880**) tasks, any of which may send an alarm, e.g., to a network management system. In this example, the predictive analytics and cross-correlations (**870**, **880**) aspects may each additionally or alternatively trigger a warning state, and notify localized processing **800** accordingly. Localized processing **800** may change the manner in which centralized processing **850** is updated accordingly.

(54) For example, when the computing system **500** is in the warning state, localized processing **800** may refrain from sending one or more of the above-discussed metrics and/or may avoid performing one or more tasks. In this way, computing resources may be conserved when there is insufficient evidence that an attack may be occurring.

(55) Although embodiments of the present disclosure require detecting an anomaly in one or more metrics before enabling further analytics, in general, it is advantageous for the metrics calculations and analytics performed to be lightweight and efficient to the extent possible. For example, during the attack warning state, the operations required in many of the examples discussed above are simple comparisons. In one such example, if the mean and modal values received by the central processing **850** functions are within a configurable window of values, an alarm may be sent that the TCP session is under attack. In this regard, values within the window may indicate that these values are nearly constant as compared to one or more previous intervals, and therefore seems to reflect malicious activity. Further, a programmable hysteresis of M intervals of time may be configured to protect against random false positives (i.e., to avoid triggering an alert for infrequent cases in which the mean and mode values are constant in the absence of an attack).

(56) As previously mentioned above, one, some, or all of the functions described above may be implemented by a computing system **500**. Such a computing system **500** may be implemented according to the hardware illustrated in FIG. 9. The example hardware of FIG. 6 comprises processing circuitry **910**, memory circuitry **920**, and interface circuitry **930**. The processing circuitry **910** is communicatively coupled to the memory circuitry **920** and the interface circuitry **930**, e.g., via one or more buses. The processing circuitry **910** may comprise one or more microprocessors, microcontrollers, hardware circuits, discrete logic circuits, hardware registers, digital signal processors (DSPs), field-programmable gate arrays (FPGAs), application-specific integrated circuits (ASICs), or a combination thereof. For example, the processing circuitry **910** may be programmable hardware capable of executing software instructions **960** stored, e.g., as a machine-readable computer program in the memory circuitry **920**. The memory circuitry **920** of the various embodiments may comprise any non-transitory machine-readable media known in the art or that may be developed, whether volatile or non-volatile, including but not limited to solid state media (e.g., SRAM, DRAM, DDRAM, ROM, PROM, EPROM, flash memory, solid state drive, etc.), removable storage devices (e.g., Secure Digital (SD) card, miniSD card, microSD card, memory stick, thumb-drive, USB flash drive, ROM cartridge, Universal Media Disc), fixed drive (e.g., magnetic hard disk drive), or the like, wholly or in any combination.

(57) The interface circuitry **930** may be a controller hub configured to control the input and output (I/O) data paths of the computing system **500**. Such I/O data paths may include data paths for exchanging signals over a communications network **105** and data paths for exchanging signals with a user. For example, the interface circuitry **930** may comprise a transceiver configured to send and receive communication signals over one or more of a cellular network, Ethernet network, or optical network. The interface circuitry **930** may also comprise one or more of a graphics adapter, display port, video bus, touchscreen, graphical processing unit (GPU), display port, Liquid Crystal Display (LCD), and Light Emitting Diode (LED) display, for presenting visual information to a user. The interface circuitry **930** may also comprise one or more of a pointing device (e.g., a mouse, stylus, touchpad, trackball, pointing stick, joystick), touchscreen, microphone for speech input, optical

sensor for optical recognition of gestures, and keyboard for text entry.

(58) The interface circuitry **930** may be implemented as a unitary physical component, or as a plurality of physical components that are contiguously or separately arranged, any of which may be communicatively coupled to any other, or may communicate with any other via the processing circuitry **910**. For example, the interface circuitry **930** may comprise output circuitry (e.g., transmitter circuitry configured to send communication signals over the communications network) and input circuitry (e.g., receiver circuitry configured to receive communication signals over the communications network). Similarly, the output circuitry may comprise a display, whereas the input circuitry may comprise a keyboard. Other examples, permutations, and arrangements of the above and their equivalents will be readily apparent to those of ordinary skill.

(59) According to embodiments of the hardware illustrated in FIG. **9**, the interface circuitry **930** is configured to receive a plurality of packets **200** of a TCP session. The processing circuitry **910** is configured to monitor for nonsequential packets of the TCP session. The nonsequential packets each comprise a sequence number **220** that is different from a next expected sequence number of the TCP session. The processing circuitry **910** is further configured to calculate, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers **220** of the nonsequential packets received in the time interval differ from the next expected sequence number. The processing circuitry **910** is further configured to determine a variation baseline representing a trend of the variation metrics over time, and detect the TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold.

(60) The present invention may, of course, be carried out in other ways than those specifically set forth herein without departing from essential characteristics of the invention. The present embodiments are to be considered in all respects as illustrative and not restrictive, and all changes coming within the meaning and equivalency range of the appended claims are intended to be embraced therein.

Claims

1. A method of detecting a Transmission Control Protocol (TCP) session hijacking attack on a TCP session, the method being implemented by a computing system and comprising: monitoring for nonsequential packets of the TCP session, the nonsequential packets each comprising a sequence number that is different from a next expected sequence number of the TCP session; calculating, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers of the nonsequential packets received in a time interval differ from the next expected sequence number; determining a variation baseline representing a trend of the variation metrics over time; detecting the TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold.
2. The method of claim 1, wherein calculating the variation metric for each of the plurality of time intervals comprises calculating: a mean difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval; or a modal difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval.
3. The method of claim 2, further comprising calculating both the mean difference and the modal difference for each of the time intervals.
4. The method of claim 3, further comprising comparing the mean difference and the modal difference of the given time interval to the mean difference and the modal difference of at least one time interval previous to the given time interval.
5. The method of claim 1, wherein: the plurality of time intervals further comprises one or more time intervals subsequent to the given time interval; detecting the TCP session hijacking attack on

the TCP session is further based on the variation metric in the subsequent time intervals being different from the variation baseline by less than the threshold.

6. The method of claim 5, wherein detecting the TCP session hijacking attack is further based on the variation metric being different from the variation baseline by less than the threshold for more than a configurable number of the time intervals.

7. The method of claim 1, further comprising commencing the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, more than a threshold number of the nonsequential packets was received.

8. The method of claim 1, further comprising commencing the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, the nonsequential packets received for the TCP session are greater in number than for one or more other TCP sessions by a threshold amount.

9. The method of claim 1, further comprising refraining from calculating the variation metric for one or more time intervals in which sequence number variation is different from the variation baseline by more than the threshold.

10. The method of claim 9, wherein the refraining from calculating the variation metric for the one or more time intervals in which the sequence number variation is different from the variation baseline by more than the threshold is responsive to the sequence number variation being different from the variation baseline for at least a configurable number of consecutive time intervals.

11. A computing system comprising: interface circuitry configured to receive a plurality of packets of a Transmission Control Protocol (TCP) session; processing circuitry communicatively coupled to the interface circuitry and configured to: monitor for nonsequential packets of the TCP session, the nonsequential packets each comprising a sequence number that is different from a next expected sequence number of the TCP session; calculate, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers of the nonsequential packets received in a time interval differ from the next expected sequence number; determine a variation baseline representing a trend of the variation metrics over time; detect a TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold.

12. The computing system of claim 11, wherein to calculate the variation metric for each of the plurality of time intervals, the processing circuitry is configured to calculate: a mean difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval; or a modal difference between the next expected sequence number and the sequence numbers of the nonsequential packets received in the time interval.

13. The computing system of claim 12, wherein the processing circuitry is further configured to calculate both the mean difference and the modal difference for each of the time intervals.

14. The computing system of claim 13, wherein the processing circuitry is further configured to compare the mean difference and the modal difference of the given time interval to the mean difference and the modal difference of at least one time interval previous to the given time interval.

15. The computing system of claim 11, wherein: the plurality of time intervals further comprises one or more time intervals subsequent to the given time interval; the processing circuitry is further configured to detect the TCP session hijacking attack on the TCP session based on the variation metric in the subsequent time intervals being different from the variation baseline by less than the threshold.

16. The computing system of claim 15, wherein the processing circuitry is further configured to detect the TCP session hijacking attack based on the variation metric being different from the variation baseline by less than the threshold for more than a configurable number of the time intervals.

17. The computing system of claim 11, wherein the processing circuitry is further configured to

commence the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, more than a threshold number of the nonsequential packets was received.

18. The computing system of claim 11, wherein the processing circuitry is further configured to commence the calculating of the variation metric for each of the plurality of time intervals in response to detecting that, in a time interval preceding the plurality of time intervals, the nonsequential packets received for the TCP session are greater in number than for one or more other TCP sessions by a threshold amount.

19. The computing system of claim 11, wherein the processing circuitry is further configured to refrain from calculating the variation metric for one or more time intervals in which sequence number variation is different from the variation baseline by more than the threshold.

20. The computing system of claim 19, wherein the processing circuitry is further configured to refrain from calculating the variation metric for the one or more time intervals in which the sequence number variation is different from the variation baseline by more than the threshold in response to the sequence number variation being different from the variation baseline for at least a configurable number of consecutive time intervals.

21. A non-transitory computer readable medium storing a computer program product for controlling a computing system, the computer program product comprising software instructions that, when run on the computing system, cause the computing system to: monitor for nonsequential packets of the TCP session, the nonsequential packets each comprising a sequence number that is different from a next expected sequence number of the TCP session; calculate, for each of a plurality of time intervals, a variation metric representing an extent to which the sequence numbers of the nonsequential packets received in a time interval differ from the next expected sequence number; determine a variation baseline representing a trend of the variation metrics over time; detect the TCP session hijacking attack on the TCP session based on the variation metric in a given time interval being different from the variation baseline by less than a threshold.
