

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388901
Kind Code	B2
Date of Patent	August 12, 2025
Inventor(s)	Tanach; Moshe et al.

Communication protocol, and a method thereof for accelerating artificial intelligence processing tasks

Abstract

A method and system for communicating artificial intelligence (AI) tasks between AI resources are provided. The method includes establishing a connection between a first AI resource and a second AI resource; encapsulating a request to process an AI task in at least one request data frame compliant with a communication protocol, wherein the at least one request data frame is encapsulated at the first AI resource; transporting the at least one request data frame over a network using a transport protocol to the second AI resource, wherein the transport protocol is different than the communication protocol; and using a credit-based flow control mechanism to transfer messages between the first AI resource and the second AI resource over the transport protocol, thereby avoiding congestion on compute resources.

Inventors: Tanach; Moshe (Bet Herut, IL), Kasus; Yossi (Haifa, IL), Khermosh; Lior (Givatayim, IL), Sivan; Udi (Zikhron Yaakov, IL)

Applicant: NeuReality Ltd. (Caesarea, IL)

Family ID: 85040566

Assignee: NeuReality Ltd. (Caesarea, IL)

Appl. No.: 18/602606

Filed: March 12, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240251016 A1	Jul. 25, 2024

Related U.S. Application Data

continuation parent-doc US 18145516 20221222 US 12132791 child-doc US 18602606
continuation parent-doc US 17387536 20210728 US 11570257 20230131 child-doc US 18145516
us-provisional-application US 63070054 20200825

Publication Classification

Int. Cl.: **H04L67/148** (20220101); **G06F9/48** (20060101); **G06F9/50** (20060101); **G06F15/173** (20060101); **H04L67/133** (20220101)

U.S. Cl.:

CPC **H04L67/148** (20130101); **G06F9/4806** (20130101); **G06F9/505** (20130101);
G06F15/17331 (20130101); **H04L67/133** (20220501);

Field of Classification Search

CPC: H04L (67/148); H04L (67/133); G06F (9/4806); G06F (9/505); G06F (15/17331)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8621627	12/2012	Eiriksson et al.	N/A	N/A
11250336	12/2021	Bernat et al.	N/A	N/A
2009/0064291	12/2008	Wahl	726/5	H04L 63/0815
2015/0012735	12/2014	Tamir et al.	N/A	N/A
2017/0286363	12/2016	Joshua et al.	N/A	N/A
2018/0288090	12/2017	Olarig et al.	N/A	N/A
2020/0089532	12/2019	Shimamura et al.	N/A	N/A
2020/0117525	12/2019	Kachare et al.	N/A	N/A
2020/0136972	12/2019	Nguyen	N/A	H04L 5/0055
2020/0219007	12/2019	Byers	N/A	G06F 9/50
2021/0011863	12/2020	Zhang	N/A	G06N 3/045
2021/0116907	12/2020	Altman	N/A	B60W 60/0015
2021/0258405	12/2020	Chourasia et al.	N/A	N/A
2021/0334542	12/2020	Shankar et al.	N/A	N/A

OTHER PUBLICATIONS

“NVM Express Over Fabrics Revision 1.1a”. NVM Express. Jul. 12, 2021. cited by applicant

Primary Examiner: Keller; Michael A

Attorney, Agent or Firm: M&B IP Analysts, LLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 18/145,516, filed Dec. 22, 2022. The Ser. No. 18/145,516 application is a continuation of Ser. No. 17/387,536, filed Jul. 28, 2021, now U.S. Pat. No. 11,570,257. The U.S. patent application Ser. No. 17/387,536 claims the benefit of U.S. Provisional Application No. 63/070,054, filed on Aug. 25, 2020, the contents of which are hereby incorporated by reference.

TECHNICAL FIELD

(1) The disclosure generally relates to communications network access and the acceleration of the processing of AI tasks within a network environment.

BACKGROUND

(2) The demand and need for efficient AI processing systems, in terms of AI computing performance, power, and cost, are increasing. These needs and demands are due in part to the increased popularity of machine learning and AI applications. The execution of such applications is performed by servers configured as a dedicated AI server or AI appliance, including software and hardware. The software may be, for example, TensorFlow®, Caffe, Pytorch®, or CNTK®, usually implementing the framework's APIs. The hardware may be, for example CPU or a combination of a CPU and a dedicated hardware accelerator, also known as a deep learning accelerator (DLA). The DLA may be, for example, GPU, ASIC, or FPGA devices.

(3) Although the DLA computation is typically implemented in hardware, the management and control of the computation is performed in software. Specifically, in an architecture that includes several dedicated hardware (HW) accelerators, there is an increased need to manage and control the jobs to be executed by the different accelerators. The management and control tasks are typically performed by an asset of software processes responsible for various functions, such as multiple tasks queue management, scheduling of jobs, drivers that interface and control the hardware programming model, etc. As such, the functionality and the performance of the entire DLA's architecture is sometimes limited by the host CPU running these processes in software.

(4) To better utilize AI compute resources in the cloud and enterprise datacenters, a disaggregation approach is being introduced. Here, primary compute resources and AI compute resources are logically and physically being disaggregated and located in separate locations in the datacenter. This allows a dynamic orchestration of the virtual machines executing AI applications on primary compute servers, as well as the AI compute resources running AI tasks on AI servers. AI tasks include, for example, machine learning, deep learning, and neural network processing tasks, for various types of applications, for example, natural language processing (NLP), voice processing, image processing, and video processing, with various usage models, for example recommendation, classification, prediction, and detection. In addition, tasks can also include preprocessing and postprocessing computation, for example, image (jpeg) decoding, non-maximum suppression (NMS) after object detection, and the like.

(5) As compute resources are disaggregated, and datacenters are being distributed, the communication between the various resources is now a performance bottleneck as it is still performed by traditional communication protocols, such as Hypertext Transfer Protocol (HTTP) over Transmission Control Protocol (TCP) or GRPC. This approach requires high CPU resources (e.g., due to the networking software stack and the networking drivers) and adds redundant latency to the processing pipeline.

(6) The traditional communication protocols are not designed to efficiently support AI computing tasks. As such, datacenters designed to support AI compute resources cannot be fully optimized to accelerate the execution of AI tasks, due to the latency and low performance of the traditional communication protocols that are not being optimized to support AI compute tasks to the clients. An optimized protocol allows to increase the efficiency of the primary/AI disaggregation in terms of latency, performance, power, and overheads as well as introducing end-to-end quality of service

features such as service level agreement (SLA) based communication, load balancing, and the like. (7) It would therefore be advantageous to provide a solution that would overcome the challenges noted above.

SUMMARY

(8) A summary of several example embodiments of the disclosure follows. This summary is provided for the convenience of the reader to provide a basic understanding of such embodiments and does not wholly define the breadth of the disclosure. This summary is not an extensive overview of all contemplated embodiments and is intended to neither identify key or critical elements of all embodiments nor to delineate the scope of any or all aspects. Its sole purpose is to present some concepts of one or more embodiments in a simplified form as a prelude to the more detailed description that is presented later. For convenience, the term “some embodiments” or “certain embodiments” may be used herein to refer to a single embodiment or multiple embodiments of the disclosure.

(9) A system of one or more computers can be configured to perform particular operations or actions by virtue of having software, firmware, hardware, or a combination of them installed on the system that in operation causes or cause the system to perform the actions. One or more computer programs can be configured to perform particular operations or actions by virtue of including instructions that, when executed by a data processing apparatus, cause the apparatus to perform the actions.

(10) In one general aspect, a method may include establishing a connection between a first AI resource and a second AI resource. The method may also include encapsulating a request to process an AI task in at least one request data frame compliant with a communication protocol, where at least one request data frame is encapsulated at the first AI resource. The method may furthermore include transporting at least one request data frame over a network using a transport protocol to the second AI resource, where the transport protocol is different than the communication protocol. The method may in addition include using a credit-based flow control mechanism to transfer messages between the first AI resource and the second AI resource over the transport protocol, thereby avoiding congestion on compute resources. Other embodiments of this aspect include corresponding computer systems, apparatus, and computer programs recorded on one or more computer storage devices, each configured to perform the actions of the methods.

(11) In one general aspect, a system may include one or more processors configured to: establish a connection between a first AI resource and a second AI resource; encapsulate a request to process an AI task in at least one request data frame compliant with a communication protocol, where the at least one request data frame is encapsulated at the first AI resource; transport the at least one request data frame over a network using a transport protocol to the second AI resource, where the transport protocol is different than the communication protocol; and use a credit-based flow control mechanism to transfer messages between the first AI resource and the second AI resource over the transport protocol, thereby avoiding congestion on compute resources. Other embodiments of this aspect include corresponding computer systems, apparatus, and computer programs recorded on one or more computer storage devices, each configured to perform the actions of the methods.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The subject matter disclosed herein is particularly pointed out and distinctly claimed in the claims at the conclusion of the specification. The foregoing and other objects, features, and advantages of the disclosure will be apparent from the following detailed description taken in conjunction with the accompanying drawings.

(2) FIG. 1 is a block diagram illustrating the communication facilitated by an AI over Fabric

(AIoF) protocol according to an embodiment.

(3) FIG. 2 is a diagram of an AIoF data frame structured by AIoF according to an embodiment.

(4) FIG. 3 is a diagram illustrating a transport of an AIoF frame over TCP according to an embodiment.

(5) FIG. 4A is a diagram illustrating a transport of an AIoF data frame over RoCE according to an embodiment.

(6) FIG. 4B is a diagram illustrating a transport of an AIoF data frame over RDMA including the AIoF basic handshaking, according to an embodiment.

(7) FIG. 5 is a flow diagram illustrating a process for establishing a connection between an AI client and an AI server according to an example embodiment.

(8) FIG. 6 is a flow diagram illustrating a dataflow between an AI client and an AI server using the AIoF protocol according to an example embodiment.

(9) FIG. 7 is a diagram of an AI acceleration system utilizing the AIoF protocol according to an embodiment.

(10) FIG. 8 is a block diagram of the AIoF transport abstraction layer according to an embodiment.

(11) FIG. 9 is a dataflow of a server-chained compute graph according to an embodiment.

DETAILED DESCRIPTION

(12) The embodiments disclosed by the invention are only examples of the many possible advantageous uses and implementations of the innovative teachings presented herein. In general, statements made in the specification of the present application do not necessarily limit any of the various claimed embodiments. Moreover, some statements may apply to some inventive features but not to others. In general, unless otherwise indicated, singular elements may be in plural and vice versa with no loss of generality. In the drawings, like numerals refer to like parts through several views.

(13) The various disclosed embodiments include a communication protocol, and method thereof allowing for high performance, low latency, and low overhead connectivity between artificial intelligence (AI) compute resources over a high-speed network fabric. The disclosed protocol further allows end-to-end performance assurance, quality of service (QoS), provision, and orchestration of the AI services. The disclosed communication protocol is referred to hereinafter as “AI over Fabric protocol” or “AIoF protocol”.

(14) The disclosed AIoF protocol enables standardized communication among several compute resources, including, a server and a client that respectively perform or respond to execution of the AI computing tasks. A server may include an AI primary compute server hosting AI applications or other applications, and the AI compute server executes AI tasks (or simply an AI task or AI job). A client may include any application or object that is utilizing the AI server for AI task offload. AI tasks include, for example, machine learning, deep learning, and neural network processing tasks, for various types of applications, for example, natural language processing (NLP), voice processing, image processing, and video processing, with various usage models, for example, recommendation, classification, prediction, and detection. In addition, tasks can also include preprocessing and postprocessing computation, for example, image (jpeg) decoding, non-maximum suppression (NMS) after object detection, and the like.

(15) The purpose of the AIoF protocol is to define an alternative communication connectivity, to a conventional processing protocol, designed to remove processing overheads and any associated latency. In an embodiment, the AIoF protocol is operable as a mediator between AI frameworks and AI computation engines. The AIoF protocol transmits and receives data frames over standard transport-layer protocols.

(16) FIG. 1 shows an example diagram **100** illustrating the communication facilitated by the AIoF protocol according to an embodiment.

(17) The AIoF protocol (schematically labeled as “**110**”) is configured to facilitate the communication between an AI client **120** and an AI server **130**. The AI client **120** is an application,

an object, and/or device utilizing the AI server **130** to offload AI tasks. The AI server **130** is an application, object, and/or device serving the AI client **120** by offloading AI task requests and responding with results. It should be noted that the AI client **120**, the AI server **130**, or both, can be realized in software, firmware, middleware, hardware, or any combination thereof.

(18) Typically, the AI client **120** would include a runtime framework **125** to execute AI applications **123**. The framework **125** may be realized using technologies including, but not limited to, TensorFlow, Caffe2, Glow, and the like, all are standardized AI frameworks or any proprietary AI framework. The AI client **120** is also configured with a set of AI APIs **127** to support standardized communication with the AI compute engine **135** at the AI server **130**.

(19) The disclosed AIoF protocol **110** is a communication protocol designed to support AI model installations and AI operations (collectively may be referred to AI computing tasks). The AIoF protocol **110** is configured to remove the overhead of a transport protocol, latency issues, and the multiple data copies required to transfer data between the AI client **120** and server **130**.

(20) In an embodiment, the AIoF protocol **110** is configured using a shared memory over the network, in which the application can use its memory while the hardware transparently copies the AI model or the data from the application memory to a network attached artificial intelligence accelerator (NA-AIA) memory via the network. As will be discussed below, the AIoF protocol provides end-to-end performance assurance and quality of service (Qos), as well as provision and orchestration of the AI services at the AI client **120**.

(21) To support the QoS, a plurality of end-to-end queues is defined for the protocol, the client, and the server to allow the level of marking to differentiate different users, flows, jobs, or queues and mark them for service priority (e.g., allowed rate, required latency, and the like). The AIoF protocol **110** includes a flow control mechanism to support multi-client multi-server topologies, that can balance traffic between multiple clients and multiple servers. The disclosed protocol further implements an end-to-end mechanism, for example, a message-based flow control or a credit-based, and the like. The flow control mechanism also allows control of the resources and provision of their compute usage, avoids congestion on the compute resources, and further allows over-provisioning of the compute resources.

(22) According to the disclosed embodiments, the AIoF protocol **110** includes a transport abstraction layer **115** configured as part of the AI client **120** and server **130**. The abstraction layer **115** is configured to fragment and de-fragment AIoF data frames, respectively, transmitted and received over a transport protocol **140**. The format of an AIoF data frame is discussed in detail below.

(23) Typically, transport protocol **140** is responsible for data integrity and retransmission in case of congestion of the link and its queues. In a further embodiment, the AIoF protocol **110** controls the integrity of the AI Job execution and contains flow control and credit information that is exchanged between the endpoints to control the scheduling and availability of AI compute resources.

(24) Different transport protocols are supported by the disclosed embodiments. The transport protocols may include a Transmission Control Protocol (TCP), a remote direct memory access (RDMA), a RDMA over converged Ethernet (ROCE), NVMe or NVMeoF, InfiniBand, and the like.

(25) The communication between the AI client **120** and AI server **130** is over a network **150**. The network **150** includes a collection of interconnected switches (not shown), allowing the connectivity between the AI client **120** and the AI server **130**. In an example configuration, the switches may include, for example, Ethernet switches. The network **150** may be a local area network (LAN), a wide area network (WAN), a metropolitan area network (MAN), and the like. The physical medium may be either a wire or a wireless medium. Typically, when deployed in datacenter, the wire medium is a copper wire or an optical fiber.

(26) The transport abstraction layers **115** of the AIoF protocol may support multiple communication channels to support the transfer of various types of data and the priority of its data.

A channel includes a separate header and control demarcations, and a separate state of operations and flow control credit related to the channel. A channel can have separate data formats and separate queues. As such, over a channel, it is possible to carry separately, in an isolated manner, a certain type of AI job traffic of that channel.

(27) The list of channels may include but is not limited to, a channel for an AI task data transfer, a channel for an AI model, a channel for control information, a channel for management, a channel for inference parameters (e.g., batch size level, required accuracy, optimization instructions/hints, unify layers, different tradeoffs), a channel for reliability and redundancy, and a channel for diagnostics and health (including, for example, a forward channel for diagnostic requests, an inference label channel to check the accuracy, and a return channel for diagnostics and health of the AI operation), and the like.

(28) The health information includes task metrics (e.g., job succeeded/failed, statistics of the results), cluster/network metrics (e.g., load on the compute, net stats, etc.), and cluster redundancy metrics. The AI metrics include supervised metrics depending on labels like accuracy results and additional non-supervised AI metrics, such as clustering of inference data, data statistics (e.g., mean, variance, histograms), and algorithm specific metrics. An example diagram illustrating the elements of the transport abstraction layer **115** at the AI client **120** is shown in FIG. **8**.

(29) The AIoF protocol **110** can support and be operational in different network topologies and be implemented in various AI acceleration systems. One example of such a system is discussed below with reference to FIG. **7**.

(30) In yet another embodiment, the AIoF protocol **110** supports a switching topology, either fabric topology like a mesh or torus or other topology or through an indirect switching topology.

(31) The supported topologies can be further utilized to transfer data over the AIoF protocol and those received at one AI server can be forwarded to another server. The specific AI server to forward the AI jobs (and data) may be designated in the AIoF data frame. The forwarding can be performed between components (e.g., CPU, AI accelerators) within the AI server. The forwarding can be performed before processing the task data in the frame's payload, according to the header information of the AIoF data frame. Alternatively, the forwarding can be performed after some level of processing of the task data that is continued in another compute server. The forwarding information is provided in the AIoF header.

(32) FIG. **2** shows an example diagram of an AIoF data frame **200** structured by AIoF according to an embodiment. The AIoF data frame **200** complies with the AIoF protocol and is being utilized to transfer data of AI tasks and results thereof. In an embodiment, AI tasks are fragmented and transferred over one or more channels supported by the AIoF protocol. In an embodiment, the frame **200** is generated and processed by a transport abstraction layer (e.g., layer **115**, FIG. **1**) of the AIoF protocol.

(33) The AIoF data frame **200** includes a header portion **210** and a payload portion **220**. The payload portion **220** is structured to carry the data to run a specific AI task. For example, the AI task may include image processing, then the data would be the image to be processed.

(34) The header portion **210** includes a number of fields designating, in part, the AI task type, the length of the payload data, a source address (or identifier), and a destination address (or identifier). The header includes the meta-data information of the AI job, including elements that are required for the processing of the AIoF frame and the AI job, channel types, information like the identifier to the job and its sources, addresses for descriptors, and job characteristics. Examples of the fields included in the header portion **210** of AIoF request frames and AIoF response frames are listed in Table 1 and Table 2, respectively.

(35) TABLE-US-00001

TABLE 1	Field Name	Description	VERSION	Version of descriptor
CHANNELS	An inference channel	(Inference, Provisioning, Management Error/Debug, Statistics,		
Inter-server	CG (Server chaining)	OPCODE	Command operation	opcode (e.g., Single Job, Batch
Job, Parallel Job, Admin)	CID	Client ID	CGID	Computational Graph (AI Mode)
ID	CREDITS	AI		

Job credits - exchanged between RX channel, allowing the AI service to get TX credits for transmission. Credit information can be returned in all AIoF messages to the client. STATUS Indication provided in response. This field can be available instead of job info.

AI_FLOW_CONTROL AI Job flow control information JID Job ID JPRIO Job priority (Strict, High, Medium, Low) QID/FLOWID (Optional) Direct targeting of Queue/Flow TIME (Optional) Time - for time sync and statistics calculations INPUTLEN Input data length JOB_FRAG_DESC Job fragmentation descriptor JINFO_AVAIL Job Information data is available in the header BATCH_SIZE Batch Size JINFOPTR Meta data pointer (Pointer to job related data, which is not actual job data) JINFOLEN Job information length JINFO_R_KEY R_KEY for JINFO data IN_SGLSIZE Input RDMA Scatter Gather List (SGL) size SGL_INADDRPTR Pointer to input data address SGL_INLEN SGL entry of input data length SGL_IN_R_KEY R_KEY for input SGL data OUT_SGLSIZE Output RDMA Scatter Gather list size SGL_OUTADDRPTR SGL pointer entry of result data address. First entry of results will be credits for following AI jobs. Credit information can be returned in all AIoF messages to the client. SGL_OUTLEN SGL entry of size of output buffer SGL_OUT_R_KEY SGL entry of R_KEY for output data

(36) TABLE-US-00002 TABLE 2 Field Name Description VERSION Version of descriptor CHANNELS An inference channel (Inference, Provisioning, Management Error/Debug, Statistics, Inter-server CG (Server Chaining)) OPCODE Command operation opcode (e.g., Single Job, Batch Job, Parallel Job, Admin) CID Client ID CGID Computational Graph (AI Mode) ID AI_FLOW_CONTROL AI Job flow control information JID Job ID JPRIO Job priority (Strict, High, Medium, Low) HDRSIZE 0 - Short Header; 1 - Long Header CREDITS AI Job credits - exchanged between RX channel, allowing the AI service to get TX credits for transmission. Credit information can be returned in all AIoF messages to the client. STATUS Indication provided in response. This field can be available instead of job info. TIME (Optional) Time - for time sync and statistics calculations OUTLEN Size of output buffer PADSIZ Byte size of padding including header size. Must be larger than 32. JOB_FRAG_DESC Job fragmentation descriptor JINFOLEN Job information length JINFO Additional job information - Statistics and additional metadata provided to the client side. RESERVED Reserved. Zero Padding for future extension. CONGESTION_NOTIFICATION Congestion notification from server to client. Higher value indicates larger congestion. Used in clients that connected to multiple servers on multiple AIoF connection to prefer less congested connection.

(37) AIoF data frame **200** is transported over a transport protocol, examples of which are provided above. When transported over a transport protocol (layer), the AIoF data frame **200** is fragmented into a number of consecutive transport layer packets, where the fragments of the AIoF frame are included in the payload portion of the transport layer packets.

(38) In an embodiment, the format of the AIoF data frame **200** can be adaptative. That is, the frame may be modified with different header fields, a header size, a payload size, and the like, or a combination thereof, to support different AI frameworks or applications. In an embodiment, the format of the data frame is negotiated, during an initialization handshake (or a discovery mode) between the AI client and server.

(39) In one configuration, several predefined formats are defined by the AIoF protocol. The version of the format can also be for a specific job or batch of jobs. In general, this flexible format can be deduced to a specific format that is selected between the two endpoints according to their capabilities, and the specific job that is currently processed.

(40) FIG. 3 shows an example diagram illustrating the transport of an AIoF data frame **310** over TCP packets **320-1** through **320-N**. As illustrated, portions of the AIoF data frame **310** are carried by the respective payload of the packets **320-1** through **320-N**. It should be noted that the size of the AIoF frame is larger than the size of the TCP packet. For example, a TCP packet's size is 100 bytes, while the size of AIoF data frame may be 1000 bytes.

(41) FIG. 4A is an example diagram for transporting an AIoF data frame **410** over ROCE packets

420-1 through **420-N** according to an embodiment. The ROCE is a network protocol that leverages RDMA to allow devices to perform direct memory to memory transfers at the application level without involving the host CPU. A standard structure of RoCE packet **420** includes a layer-4 packet header (UDP), an ROCE header, and a RoCE payload. The AIOF data frame **410** is first encapsulated in a RDMA frame **430** and then into consecutive ROCE packets **420-1** through **420-N**.

(42) As illustrated in FIG. 4A, portions of the AIOF data frame **410** are carried out by the respective payload of ROCE packets **420-1** through **420-N**. It should be noted that the size of the AIOF frame is larger than the size of the ROCE packet.

(43) FIG. 4B is an example diagram of transporting an AIOF data frame **440-1** through **440-N** over ROCE packets **460-1** through **460-N**, following the AIOF handshake, according to an embodiment. An example diagram illustrating an AIOF handshake is shown in FIG. 6. In an example embodiment, the AIOF data frames **440-1** through **440-N** are encapsulated in RDMA frames **450-1** through **450-N**, more particularly, with specific commands such as SEND and READ in the payload of each packet. Portions of the AIOF frame are carried out by corresponding payloads of the ROCE packets **460-1** through **460-N**. In an embodiment, the payload is read from the client using RDMA read operation that may include, but not limited to, read, read response, and the like. It should be noted that the AIOF frame header can be sent separately from the AIOF job data itself.

(44) FIG. 5 is an example flow diagram illustrating a method for establishing connection between an AI client **120** and an AI server **130** according to an embodiment. It should be noted that all steps may be optional and may be performed offline to enable the link to start with a pre-shared configuration.

(45) At **S501**, a connection is initiated by the AI client **120**, which sends a list of provision requests for a new connection. The list of provisions may include, but is not limited to, a client ID, a computational graph service level agreement (CG_SLA), and a computational graph (CG) descriptor. The AI server **130** receives the list, and client connection provisioning occurs in the hardware. At **S502**, a response is sent by the AI server **130**. The response may indicate the success or failure of the connection.

(46) At **S503**, an AIOF administrator (Admin) channel creation is requested. Such a channel may be used for the initiation of the AIOF and transport protocol (e.g., RDMA) connections. The Admin channel may further regulate query and response messages for management and status updates such as, but not limited to, status and statistic gathering, state changes, and event alerts. In an embodiment, the Admin channel may resize on an RDMA and/or TCP. At **S504**, an administrator channel completion information is sent from the AI server **130** to the AI client **120**.

(47) At **S505**, the transport connection request is sent from the AI client **120** to the AI server **130**. At **S506**, the connection completion information is sent from the AI server **130** to the AI client **120**.

(48) At **S507**, an AIOF connection message is sent from the AI client **120** to the AI server **130**. Such connection message includes transient AIOF link connection information, but is not limited to, a client ID and computational graph ID (CG_ID). A network connection is configured at the AI server **130** for mapping between a queue pair (QP), an input queue, a flow ID, Job_ID, credits, and AI Job Scatter Gather List (SGL) parameters. The Job ID is used for initialization and the credits are allocated for AIOF flow control. At **S508**, a response message is sent to the AI client **120** indicating the success or failure of the AIOF connection establishment.

(49) FIG. 6 illustrates an example flow diagram illustrating the dataflow between an AI client **120** and an AI server **130** using the AIOF protocol to transport AI computing tasks according to an embodiment. In the example shown in FIG. 6, the transport is over an RDMA serving as a transport layer. The AIOF and RDMA components at the client **120** side are labeled as **610** and **620**, while the AIOF and RDMA components at the server **130** side are labeled as **630** and **640**, respectively.

(50) It should be noted that the AIOF and RDMA components may include software stack and/or circuits configured to execute the various tasks related to the operation of the AIOF and RDMA

protocol. The AIOF component (either **610** or **630**) implements the abstraction layer (e.g., layer **115**, FIG. **1**).

(51) The RDMA protocol provides the ability to access (read, write) memory on a remote system (e.g., AI client or server) without interrupting the processing of the CPUs on that system. There are a number of transfer commands in RDMA including SEND—a request to send data; ACK—acknowledgment of a SEND request, WRITE—write data into a remote (virtual) memory, and READ—read data out of the remote (virtual) memory. These commands are used when the AIOF is carried over RDMA/ROCE. The flow describes the operation of the AIOF and it is in addition to the regular RDMA/ROCE handshake for initialization and regular operation.

(52) In an embodiment, when using TCP/IP packets, the AIOF data frames are carried over Ethernet SEND/RECEIVE packets, encapsulated over TCP/IP, in addition to the regular TCP/IP message protocols. In this embodiment, the handshake can also be implemented over layer-7 protocols, such as HTTP or HTTP2, where the messages will be encapsulated over the framing of these protocols.

(53) At **S601**, an AIOF frame is sent from the AI client **120** to the AI server **130**. The AIOF frame is transported over the network using an RDMA SEND command. The AIOF frame (e.g., frame **200**, FIG. **2**) designates at least a client ID (CID), a memory address (or pointer) to a client memory where to read the task data from, a job (or task) ID (JID), and the task data to be processed. This is a request from the client to initiate a job which is afterwards followed by the RDMA READ command from the AI server **130**. Alternatively, TX credits can be used, where such credits guarantee the client the ability to transmit the job if credit is available. This saves round-trip time (RTT) of the link latency for the job.

(54) At **S602**, the AIOF component **630** sends an AIOF data frame (“read job data”) including the job ID (JID) of the task requested by the client, and the client ID (CID). This AIOF data frame is translated to an RDMA READ request sent to the RDMA component **620**, to directly read the task data from the client memory. At **S603**, the read task data is sent back to the server-side AIOF component **630**. For TCP/IP messages will be carried with SEND frames.

(55) At **S604**, when the processing of the AI task is completed by the AI server **130**, another AIOF data frame (“done”) is sent to the client side AIOF component **610** from the RDMA component **640**. The frame is transferred by means of the RDMA SEND command. In an embodiment, such a frame is configured to designate the client ID (CID) and job ID (JID). For TCP/IP messages will be carried with SEND frames.

(56) At **S605**, an AIOF data frame (“send result data”) is issued by the server side AIOF component **630**, such a command is translated to an RDMA SEND command to provide results to the AI client **120**. In another embodiment, an AIOF data frame (“write result”) is issued by the server-side AIOF component **630**, such a command is translated to an RDMA WRITE to write the result to the client's memory, indicated by the memory pointer (or address). RDMA WRITE may be an immediate WRITE to indicate the occurrence and completion of the WRITE operation to the AI client side **120**. In an embodiment, TCP/IP messages are carried with SEND frames. If data is transferred by a SEND command, then data is copied to the designated address space afterward at the client side **120**.

(57) In an embodiment, the “read results” and “write results” commands are tracked using a table that records the client ID (CID), and an address (pointer) for the client's memory to write the results. Upon, a write request command, the address of the memory is retrieved from the table based on the client ID and job ID. If the address is invalid, an error message is returned. In an embodiment, the results are sent to the AI client **120** using an RDMA SEND operation, thus the tracking table is not required.

(58) It should be noted that **S605** and **S606** may be iteratively performed multiple times until the entire results data are stored in the memory.

(59) At **S606**, an RDMA ACK message is sent to the AI server **130**, acknowledging the completion of the RDMA SEND and/or WRITE. For TCP/IP, a message will be carried with SEND frames.

(60) In an embodiment, where the SEND and WRITE messages are used to transfer the results, the AI client **120** is aware of process completion without any further messages. In an optional embodiment at **S607**, an AIOF data frame (“done”) indicating the completion of the transaction, is sent to the client side AIOF component **610** from the AIOF component **630**. In this example, the frame is transferred using an RDMA SEND. For TCP/IP, the message will be carried with SEND frames.

(61) The AIOF components may be implemented in software, firmware, middleware, hardware, or any configuration thereof. Further, such components can be realized as virtual software including entities software containers, virtual machines, microservices, and the like. In an embodiment, the AIOF component can be integrated into a network interface card (NIC) included in the server or client. Such integration can be achieved using a protocol software or firmware driver.

(62) It should be noted that the disclosed embodiments are not limited for the transport of an AI service over ROCE only, but the AIOF protocol can utilize any transport protocol for AI services and/or other types of services. For example, a transport layer or protocol may include the TCP/IP.

(63) A security layer is particularly important in a disaggregated system as disclosed where data flows between different AI resources and clients, at different security levels. To this end, it should be noted that the AIOF can optionally reside in a secured protocol that authenticates the endpoints of the connection (e.g., client and server) while keeping the data confidential by encrypting the transferred data. Such a configuration incorporates a strong encryption mechanism of the protocol to avoid security attacks, such as man-in-the-middle attacks, eavesdropping, and data replication. In an embodiment, security may be implemented at the transport protocol level. For example, IPsec may be implemented at an IP level that is relevant for UDP and TCP transport that carries AIOF over RDMA and AIOF over TCP, respectively. In another example, transport layer security (TLS) may be implemented as an end-to-end client-to-server, security protocol for TCP-based transport. Security can also be implemented in the AIOF layer as part of the AIOF protocol while securing the AIOF payload according to the security indicators compounded from the AIOF header information. The Security Association can be achieved with the AI application level identifiers, for example, the CG ID, the Client_ID, channel, or the AI job identifiers and characteristics.

(64) In an embodiment, the AIOF protocol may be transported by transport protocols with a strong encryption. The AIOF entities may be mapped into the transport protocol entities for encryption so that the AI client or server is identified, for example, by an IP endpoint or by a 5-tuple network ID for the IPSEC and TLS case. In an example embodiment, the AIOF information for an AI task including, without limitation, a computer graph (model) of the user, data used for inference, and response data from the server, are all encrypted and authenticated. Furthermore, each AI client is provided with a separate connection and security association (SA) that ensures isolated encryption channels. In this scenario, a virtual path at the server side and the Artificial Intelligence Software Solutions (AISS) are configured to ensure secured AI computing at the server for each specific client. Such configuration achieves a condition similar to a physical allocation of resources at client infrastructures for higher security. This continues the secure isolation provided through the AIOF protocol.

(65) FIG. 7 shows an example diagram of an AI acceleration system **700** utilizing the AIOF protocol. The system **700** is utilized to accelerate the operations of AI processing tasks including, but not limited to, deep-learning tasks.

(66) The system **700** includes a plurality of application servers **710-1** through **710-N**, at least one appliance server **720**, and a number of switches **730** connecting the various components of the system **700**. In an embodiment, each of the switches **730** is an Ethernet switch.

(67) Each application server **710** executes an AI application over an AI software framework. Such a framework may be, for example, TensorFlow®, Caffe, Pytorch®, or CNTK®, and the like. Other frameworks include an application service working as an HTTP client with a remote AI server, such as a Tensorflow® Serving, NVIDIA® Triton server, and Pytorch® serving. In an example

embodiment, the application server **710** may include a central processing unit (CPU) **711**, a network interface card (NIC) **712**, and a memory **713**.

(68) Each application server **710** further communicates with the artificial intelligence accelerator (AIA) appliance server **720** that executes specific AI processing tasks. The AIA appliance server **720** is designed according to the disclosed embodiments to accelerate the execution of the AI tasks. The AIA appliance server **720** may include, in one configuration, an AIA switch **721** with a plurality of NA-AIAs **722**. Each NA-AIA **722** is coupled to its own memory **722-M**. The AIA switch **721** may be an Ethernet switch. The AIA appliance server **720** may be viewed as the AI server (**130**, FIG. **1**).

(69) The system **700** provides a scalable solution as there are no compute and communication bottlenecks. Specifically, in an embodiment, additional NA-AIAs can be connected to the switch **721**, thereby expanding the switch **721** to the AIA appliance server **720**. The switch **721** is configured to have enough ports and bandwidth to allow data traffic to be transferred flawlessly between the application servers **710** and the AIA appliance server **720**.

(70) According to the disclosed embodiments, further acceleration is achieved by utilizing the disclosed AIoF protocol. Specifically, an acceleration is achieved based on the AIoF protocol to write and/or read to and from a shared memory over network. That is, an AI application can use its memory while the hardware transparently copies the AI model or the data from the application memory to an NA-AIA's memory via the network.

(71) It should be noted that the AIoF protocol can be implemented in other topologies of AI acceleration systems or direct-attached acceleration systems while providing the described benefits of AI end-to-end QoS and efficient AI job framing and transmission. The AIoF protocol can be implemented with systems including ROCE/TCP and the protocol can run in software on the server side. The AIoF protocol can be implemented with systems that utilize general-purpose CPUs for AI tasks and dedicated hardware.

(72) FIG. **8** illustrates an example diagram of the transport abstraction layer **115** at the AI client (**120**, FIG. **1**). The element includes a runtime engine **810**, a provisioning engine **820**, and an orchestration engine **830**. The runtime engine **810** syncs all the operations of the requests received from the AI Application. The provisioning engine **820** provides the programming interface to the provisioning functionality. The orchestration engine **830** provides the programming interface to the management and orchestrate functionality.

(73) According to the disclosed embodiments, the disclosed protocol can support a server chaining. Server chaining functionality allows the spanning of a single compute graph (CG) over multiple AI servers. Due to the overhead of server chaining functionality, it is recommended to strive to separate complete compute graphs into different AI servers, if possible, and use a server chaining function as a backup option. Note, that separating a complete compute graph can also be in the same server and not necessarily be implemented with server chaining.

(74) Here, an AI client (e.g., client **120**, FIG. **1**) maintains a single AIoF connection for a server-chained compute graph. That is, the AI client receives the responses from the same AI server that the client sends the requests to. As such, server-chaining of compute graphs is transparent to clients.

(75) Also required by the protocol is that inter-server requests and responses use the same AIoF connection between the pair of servers. That is a server that forwards the request to the next server, receives the response from this server on the same AIoF connection.

(76) FIG. **9** illustrates dataflow of a server-chained compute graph according to an embodiment.

(77) The AI client **910** sends a job request to and receives a job response from a first AI server **920**. The job includes the compute graph to process using server chaining. The connection between the AI client **910** and the AI server **920** is over an AIoF connection **901**. The first AI server **920** establishes another single connection **902** with one or more AI servers (collectively marked as **930**). The control and data are transferred over the AIoF connections as discussed above.

(78) It should be noted that AI server 920 is configured to initiate AIoF connection to a distant server (during compute graph provisioning), build and send job requests, receive job responses, receive 'JobReadDone' and release buffer, respond to RDMA READ requests, and SEND and/or WRITE the receiver of the JobResponse.

(79) The embodiments disclosed herein can be implemented as hardware, firmware, software, or any combination thereof. The application program may be uploaded to, and executed by, a machine comprising any suitable architecture. Preferably, the machine is implemented on a computer platform having hardware such as one or more central processing units ("CPUs"), a memory, and input/output interfaces.

(80) The computer platform may also include an operating system and microinstruction code. The various processes and functions described herein may be either part of the microinstruction code or part of the application program, or any combination thereof, which may be executed by a CPU, whether or not such computer or processor is explicitly shown.

(81) In addition, various other peripheral units may be connected to the computer platform such as an additional network fabric, storage unit, and a printing unit. Furthermore, a non-transitory computer readable medium is any computer readable medium except for a transitory propagating signal.

(82) It should be understood that any reference to an element herein using a designation such as "first," "second," and so forth does not generally limit the quantity or order of those elements. Rather, these designations are generally used herein as a convenient method of distinguishing between two or more elements or instances of an element. Thus, a reference to first and second elements does not mean that only two elements may be employed there or that the first element must precede the second element in some manner. Also, unless stated otherwise, a set of elements comprises one or more elements.

(83) As used herein, the phrase "at least one of" followed by a listing of items means that any of the listed items can be utilized individually, or any combination of two or more of the listed items can be utilized. For example, if a system is described as including "at least one of A, B, and C," the system can include A alone; B alone; C alone; A and B in combination; B and C in combination; A and C in combination; or A, B, and C in combination.

(84) All examples and conditional language recited herein are intended for pedagogical purposes to aid the reader in understanding the principles of the disclosure and the concepts contributed by the inventor to furthering the art, and are to be construed as being without limitation to such specifically recited examples and conditions.

Claims

1. A method for communicating artificial intelligence (AI) tasks between AI resources, comprising: establishing a connection between a first AI resource and a second AI resource; encapsulating a request to process an AI task in at least one request data frame compliant with a communication protocol, wherein the at least one request data frame is encapsulated at the first AI resource; transporting the at least one request data frame over a network using a transport protocol to the second AI resource, wherein the transport protocol is different than the communication protocol; and using a credit-based flow control mechanism to transfer messages between the first AI resource and the second AI resource over the transport protocol, thereby avoiding congestion on compute resources; wherein the first AI resource is an AI client and the second AI resource is at least one AI server.

2. The method of claim 1, wherein the transport protocol provisions transport characteristics of the AI task.

3. The method of claim 1, further comprising: transferring using the credit-based flow control mechanism credit information, wherein the credit information controls scheduling and availability

of the compute resources.

4. The method of claim 3, wherein the credit information is included in a header portion of a message of the transport protocol.

5. The method of claim 4, wherein the credit information includes a number of credits allocated per AI job, wherein the AI task includes one or more AI jobs.

6. The method of claim 5, wherein the credits are allocated to the client AI.

7. The method of claim 2, wherein the at least one AI server is configured to accelerate execution of the AI task.

8. The method of claim 1, wherein the transport protocol is any one of: a Transmission Control Protocol (TCP), a remote direct memory access (RDMA), a RDMA over converged Ethernet (RoCE), NVMe or NVMeoF, and an InfiniBand.

9. A system for communicating artificial intelligence (AI) tasks between AI resources comprising: one or more processors configured to: establish a connection between a first AI resource and a second AI resource; encapsulate a request to process an AI task in at least one request data frame compliant with a communication protocol, wherein the at least one request data frame is encapsulated at the first AI resource; transport the at least one request data frame over a network using a transport protocol to the second AI resource, wherein the transport protocol is different than the communication protocol; and use a credit-based flow control mechanism to transfer messages between the first AI resource and the second AI resource over the transport protocol, thereby avoiding congestion on compute resources; wherein the first AI resource is an AI client and the second AI resource is at least one AI server.

10. The system of claim 9, wherein the transport protocol provisions transport characteristics of the AI task.

11. The system of claim 10, wherein at least one AI server configured to accelerate execution of the AI task.

12. The system of claim 9, wherein the one or more processors are further configured to: transfer using the credit-based flow control mechanism credit information, wherein the credit information controls scheduling and availability of the compute resources.

13. The system of claim 12, wherein the credit information is included in a header portion of a message of the transport protocol.

14. The system of claim 13, wherein the credit information includes a number of credits allocated per AI job, the AI task includes one or more AI jobs.

15. The system of claim 14, wherein the first AI resource is a client AI, and the credits are allocated to the client AI.

16. The system of claim 9, wherein the transport protocol is any one of: a transmission control protocol (TCP), a remote direct memory access (RDMA), a RDMA over converged Ethernet (RoCE), NVMe or NVMeoF, and an InfiniBand.

17. A non-transitory computer-readable medium storing a set of instructions for communicating artificial intelligence (AI) tasks between AI resources, the set of instructions comprising: one or more instructions that, when executed by one or more processors of a device, cause the device to: establish a connection between a first AI resource and a second AI resource; encapsulate a request to process an AI task in at least one request data frame compliant with a communication protocol, wherein the at least one request data frame is encapsulated at the first AI resource; transport the at least one request data frame over a network using a transport protocol to the second AI resource, wherein the transport protocol is different than the communication protocol; and use a credit-based flow control mechanism to transfer messages between the first AI resource and the second AI resource over the transport protocol, thereby avoiding congestion on compute resources; wherein the first AI resource is an AI client and the second AI resource is at least one AI server.
