

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12388915
Kind Code	B1
Date of Patent	August 12, 2025
Inventor(s)	Petersen; Brian Arnold et al.

Protocol independent data unit forwarding

Abstract

Disclosed are systems, methods, and non-transitory computer-readable media for protocol independent data unit forwarding. A packet forwarding system receives a data unit comprising a header byte string via an input port. The packet forwarding system parses the data unit based on a header type determined based on the input port, yielding a parsing output describing the header byte string of the data unit. The packet forwarding system updates a metadata item associated with the data unit based on the parsing output and determines a packet forwarding instruction for forwarding the data unit to a destination based on the metadata item associated with the data unit. The packet forwarding system forwards the data unit to the destination based on the packet forwarding instruction and the metadata item associated with the data unit.

Inventors: Petersen; Brian Arnold (San Francisco, CA), Myers, Jr.; Roy T. (Morgan City, CA), Engelkemier; Darren S. (Menlo Park, CA), Sedarat; Hossein (San Jose, CA), Shirani; Ramin (Morgan Hill, CA), Wellbaum; Tom Quoc (San Jose, CA), Leung; Poon-Kuen (Fremont, CA)

Applicant: Ethernovia Inc. (San Jose, CA)

Family ID: 82385191

Assignee: Ethernovia Inc. (San Jose, CA)

Appl. No.: 18/468908

Filed: September 18, 2023

Related U.S. Application Data

continuation parent-doc US 17807255 20220616 US 11818236 child-doc US 18468908
continuation parent-doc US 16926057 20200710 US 11394813 child-doc US 17807255
us-provisional-application US 63013440 20200421
us-provisional-application US 62872458 20190710

Publication Classification

Int. Cl.: **H04L69/18** (20220101); **H04L49/00** (20220101); **H04L49/25** (20220101); **H04L69/22** (20220101)

U.S. Cl.:

CPC **H04L69/18** (20130101); **H04L49/252** (20130101); **H04L49/3027** (20130101);
H04L69/22 (20130101); H04L2212/00 (20130101)

Field of Classification Search

CPC: H04L (69/18); H04L (49/252); H04L (49/3027); H04L (69/22); H04L (2212/00)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7558270	12/2008	Wilford	370/395.42	H04L 45/583
8793358	12/2013	Godbole	709/223	H04L 45/40
10778588	12/2019	Singh	N/A	N/A
11394813	12/2021	Petersen et al.	N/A	N/A
2007/0019661	12/2006	Rowett et al.	N/A	N/A
2015/0281106	12/2014	Lee	370/236	H04L 45/38
2016/0357781	12/2015	Bastide et al.	N/A	N/A
2017/0279638	12/2016	Albrecht	N/A	H04L 69/22
2018/0097722	12/2017	Callard	N/A	H04L 12/4641
2019/0289102	12/2018	Goel	N/A	H04L 45/74591
2021/0144405	12/2020	Sharangpani	N/A	N/A
2021/0160350	12/2020	Volpe	N/A	H04L 67/563

OTHER PUBLICATIONS

U.S. Appl. No. 16/926,057 U.S. Pat. No. 11,394,813, filed Jul. 10, 2020, Protocol Independent Data Unit Forwarding. cited by applicant

U.S. Appl. No. 17/807,255, filed Jun. 16, 2022, Protocol Independent Data Unit Forwarding. cited by applicant

“U.S. Appl. No. 16/926,057, Non Final Office Action mailed Oct. 15, 2021”, 13 pgs. cited by applicant

“U.S. Appl. No. 16/926,057, Notice of Allowance mailed Mar. 17, 2022”, 10 pgs. cited by applicant

“U.S. Appl. No. 16/926,057, Response filed Jan. 18, 2022 to Non Final Office Action mailed Oct. 15, 2021”, 11 pgs. cited by applicant

“U.S. Appl. No. 17/807,255, Non Final Office Action mailed Feb. 15, 2023”, 23 pgs. cited by applicant

“U.S. Appl. No. 17/807,255, Notice of Allowance mailed Jul. 12, 2023”, 5 pgs. cited by applicant

“U.S. Appl. No. 17/807,255, Preliminary Amendment filed Oct. 28, 2022”, 8 pgs. cited by applicant

“U.S. Appl. No. 17/807,255, Response filed May 25, 2023 to Non Final Office Action mailed Feb. 15, 2023”, 11 pgs. cited by applicant

Primary Examiner: Hailu; Kibrom T

Attorney, Agent or Firm: Schwegman, Lundberg & Woessner, P.A.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application is a continuation of U.S. patent application Ser. No. 17/807,255, filed on Jun. 16, 2022, which is a continuation of U.S. patent application Ser. No. 16/926,057, filed on Jul. 10, 2020, which claims the benefit of priority of U.S. Provisional Application No. 62/872,458, filed on Jul. 10, 2019, and

TECHNICAL FIELD

(1) An embodiment of the present subject matter relates generally to data forwarding, and more specifically, to protocol independent data unit forwarding.

BACKGROUND

(2) Computer networking allows for the exchange of data between various interconnected devices. Data transmitted from one device to another is broken into smaller data units (e.g., data packet, data frames), which are each transmitted separately along a network path to their destination. Each data unit consists of an ordered set of individual values, such as a binary number made up of an ordered set of zeroes and ones (e.g., 0001, 0010, etc.) or contiguous groups of eight binary digits (bytes). The data units are generated using various communication protocols (e.g., Ethernet, Internet Protocol, (IP), User Datagram Protocol (UDP)) that define the rules, syntax and semantics for arranging a data unit. For example, a communication protocol may define a length of the data unit (e.g., number of individual values (bytes)), as well as identify the types of data that are to be stored in the various positions of the data unit, such as identifying the range of bytes (e.g., bytes 0-127) that store the header, transferred data (e.g., payload), and/or trailer.

(3) The header may include data used to deliver a data unit to its intended destination, such as by including data identifying the data unit type (e.g., the communication protocol used to generate the data unit), a source address, a destination address, and the like. Multiple communication protocols may be used to transmit data through use of data encapsulation. Encapsulation is a method in which a data unit generated using one communication protocol is stored as the payload of a data unit generated using a different communication protocol or another instance of the same protocol. The header type of the outermost header defines the data unit type of the data unit. In some cases, a data unit may be the encapsulated multiple times resulting in a data unit that is encapsulated in multiple layers of headers.

(4) Current data forwarding systems provide limited functionality when processing encapsulated data units. For example, current data forwarding systems do not modify or replace the innermost header, resulting in restrictions on the types of devices that are able to communicate directly via a network. A common work-around for such restrictions is to deliver a data unit to a Central Processing Unit (CPU), where software is used to re-cast the data unit into a different data unit type. While this approach can work, it presents significant performance bottlenecks and greatly increases the resource and power consumption of the system and network.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Various ones of the appended drawings merely illustrate example embodiments of the present disclosure and cannot be considered as limiting its scope.

(2) FIG. 1 is a block diagram of a protocol independent data unit forwarding system, according to certain example embodiments.

(3) FIG. 2 is a block diagram showing a processing sequence performed by the parser **104**, according to some example embodiments.

(4) FIG. 3 is a flowchart showing a method for receiving and initial processing of a data unit for protocol independent data unit forwarding, according to some example embodiments.

(5) FIG. 4 is a flowchart showing a method for parsing a data unit for protocol independent data unit forwarding, according to some example embodiments.

(6) FIG. 5 is a flowchart showing a method of performing a header lookup for protocol independent data unit forwarding, according to some example embodiments.

(7) FIG. 6 is a flowchart showing a method of protocol independent data unit forwarding, according to some example embodiments.

(8) FIG. 7 is a flowchart showing a method of forwarding a data unit to its destination based on a forwarding instruction, according to some example embodiments.

(9) FIG. 8 is a block diagram illustrating an example software architecture, which may be used in conjunction with various hardware architectures herein described.

(10) FIG. 9 is a block diagram illustrating components of a machine, according to some example embodiments, able to read instructions from a machine-readable medium (e.g., a machine-readable storage medium) and perform any one or more of the methodologies discussed herein.

DETAILED DESCRIPTION

(11) In the following description, for purposes of explanation, various details are set forth in order to provide a thorough understanding of some example embodiments. It will be apparent, however, to one skilled in the art, that the present subject matter may be practiced without these specific details, or with slight alterations.

(12) Disclosed are systems, methods, and non-transitory computer-readable media for protocol independent data unit forwarding. A protocol independent data unit forwarding system is configured to process and forward data units of various communication protocol types, such as Ethernet, IP, IPv4, IPv6, UDP, Audio Video Transport Protocol (AVTP), and the like. The protocol independent data unit forwarding system includes multiple functional components providing functionality to properly process and forward a data unit to its intended destination. For example, the functional components may provide functionality such as interpreting headers included in a data unit, terminating tunnels if needed, identifying a data unit's destination, re-encapsulating the payload as needed, and the like. The functional components of the data unit forwarding system are communication protocol type agnostic, meaning that the functional components use a common set of operations, logic, templates, and/or tables to process a data unit regardless of its communication protocol type. This is in contrast to current forwarding systems that employ specialized sets of functions for different communication protocol types. As a result, the protocol independent data unit forwarding system is not limited in the types of data units it can process or the types of network segments to which a data unit may be forwarded.

(13) FIG. 1 is a block diagram of a protocol independent data unit forwarding system **100**, according to certain example embodiments. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components (e.g., modules, mechanisms, devices, nodes, etc.) that are not germane to conveying an understanding of the inventive subject matter have been omitted from FIG. 1. However, a skilled artisan will readily recognize that various additional functional components may be supported by the protocol independent data unit forwarding system **100** to facilitate additional functionality that is not specifically described herein.

(14) The protocol independent data unit forwarding system **100** provides for the exchange of data between various interconnected devices. The devices may be of any of a variety of types of devices capable of network communication with other devices. For example, each device can include some or all of the features, components, and peripherals of the machine **900** shown in FIG. 9. Similarly, the protocol independent data unit forwarding system **100** may be implemented on any of a variety of types of devices capable of network communication with other devices, such as a device including some or all of the features, components, and peripherals of the machine **900** shown in FIG. 9.

(15) As shown, the protocol independent data unit forwarding system **100** includes multiple functional components connected in series to provide a data processing pipeline for protocol independent data unit forwarding. Each functional component may be a physical data processing element configured to perform various data processing functions based on received input data. A functional component may generate a data output, which may be provided as input into other functional components in the data processing pipeline.

(16) As shown, the protocol independent data unit forwarding system **100** includes a receive port manager **102**, a parser **104**, a header lookup **106**, a lookup result processor **108**, a data buffer **110**, a header rewrite **112**, and a transmit port manager **114**.

(17) The receive port manager **102** embodies various Media Access Controllers (MACs) and operates to receive data units transmitted in a network via an assortment of external ports. The external ports may include Ethernet, as well as a variety of specialty ports such as Camera Serial Interface (CSI), Controller Area Network (CAN), Local Interconnect Network (LIN), and the like. Each port is identified by a unique port identifier.

(18) After receiving a data unit, the receive port manager **102** may perform a variety of data processing operations to generate an output that can be provided to the other functional components of the protocol independent data unit forwarding system **100**. For example, the receive port manager **102** may initially perform any media specific adaptations on the received data unit as well as multiplex the data units of the several receive ports to a multiplexed path of the data processing pipeline such that the data unit, or any portion, thereof may be transmitted to the other downstream functional components of the protocol independent data unit forwarding system **100**.

(19) After multiplexing the data unit, the receive port manager **102** identifies a header byte string from the data unit. The header byte string includes a subset of the ordered set of individual values (e.g., bytes) from the data unit. For example, the header byte string may include a predetermined number of the first ordered values (e.g., first *n* bytes) from the data unit. The predetermined number may be any number that is less than or equal to the total number of values in the data unit, however the predetermined number used may be selected to encompass all relevant header data included in the data unit, while remaining small enough that the size of the header byte string does not provide undue burden during subsequent processing of the header byte string.

(20) The receive port manager **102** may also generate a metadata item for the data unit that describes the context of the data unit. For example, the metadata item includes a series of values that describe a context of the data unit that can be used by the various functional components of the protocol independent data unit forwarding system **100** to process the data unit. For example, a functional component may use the values in the metadata item to determine which operation the functional component is to perform to further process the data unit. As each individual operation is performed, the metadata item may be updated to further define the context of the data unit. For example, the metadata item may be updated based on the resulting output of an operation by adding new values to the metadata item, modifying existing values in the metadata item, and the like.

(21) The receive port manager **102** configures some initial data values in the metadata item. For example, the receive port manager **102** may include data values in the metadata item to identify the port via which the data unit was received (e.g., unique port identifier), a data length of the data unit (e.g., byte count), and/or data identifying any errors that may have been detected during reception of the data unit.

(22) As shown, the receive port manager **102** may provide data to the parser **104** and/or the data buffer **110** via the data processing pipeline. The parser **104**, header lookup **106**, and lookup result processor **108** are configured to process a data unit based on the header byte string and the metadata item. That is, the operations performed by the parser **104**, header lookup **106**, and lookup result processor **108** can be performed without any portion of the data unit following the header byte string.

(23) Accordingly, the receive port manager **102** may provide the parser **104** with only the header byte string and the metadata item. In this type of embodiment, the receive port manager **102** may copy the portion of the data unit proceeding the header byte string to the data buffer **110** to bypass the parser **104**, header lookup **106**, and lookup result processor **108**. As a result, the portion of the data unit following the header byte string is not transmitted through parser **104**, header lookup **106**, and lookup result processor **108**, which reduces overall resource usage associated with processing the data unit.

(24) Providing the parser **104** with only the header byte string and the metadata item is only one embodiment, however, and is not meant to be limiting. In some embodiments, the receive port manager **102** may transmit the metadata item and the entire data unit or a portion of the data unit that is greater than the header byte string to the parser **104**. Further, in some embodiments, the receive port manager **102** may transmit the entire data unit or the portion of the data unit following the header byte string to the data buffer **110** along with some portion of the header byte string to the data buffer **110**. Alternatively, the receive port manager **102** may not transmit any data to the data buffer **110**, such as in embodiments in which the receive port manager **102** provides the entire data unit to the parser **104**. In this type of embodiment, the protocol independent data unit forwarding system **100** may not include the bypass path to the data buffer **110**.

(25) The parser **104** is configured to determine the types and locations of the headers that are present in the data unit (e.g., in the header byte string). For example, the parser **104** determines the types of the headers and the offset and width (e.g., a range pair) defining the location of the headers in the data unit.

(26) In general, headers do not include encoded values that identify the header's type. Rather, a previous (e.g., outer) header indicates the header type of the next sequential header in the data unit. Accordingly, the parser **104** uses a header-by-header incremental process to determine the overall structure of the headers included in a data unit.

(27) As the outermost (e.g., first) header does not have a previous header from which to access a next header type, the parser **104** uses the metadata item received from the receive port manager **102** to identify the header type of the outermost (e.g., first) header. Specifically, the parser **104** identifies the header type of the outermost header based on the data identifying the port via which the data unit was received (e.g., unique port identifier). For example, each port or port identifier may be mapped to a header type. Accordingly, the parser **104** determines the header type mapped to the unique port identifier received in the metadata items, which is inferred to be the header type of the outermost header. The mapping between the ports and the header types of the outermost header allows the parser **104** to operate with any of a variety of network types.

(28) The parser **104** parses a layer of headers prior to providing an output to the header lookup **106**. A layer of headers is generally a collection of one or more headers that are closely related in a manner such that information from each of the headers is used to analyze or understand the data included in the collection headers. For example, an Ethernet header and a Virtual Local Area Network (VLAN) header can be grouped as a single layer of headers because data included in the VLAN header is used to provide context for the Ethernet header. Similarly, an IPv4 or IPv6 header that is followed by a UDP header may be grouped as a single layer because the UDP header may act as an extension of an IPv4 or IPv6 header with regards to identifying the header type that follows the UDP header. The definition of a header layer and the functionality of parser **104** may be configured as desired such that a header layer does not necessarily have to include headers that are related in the manner described above. For example, the functionality of the parser **104** may be configured such that a header layer includes unrelated headers to account for circumstances where it may be ideal to do so. Similarly, the parser **104** may be configured to split related headers into separate header layers.

(29) Rather than extract data fields from the headers to provide to downstream functional components as separate values, the parser **104** generates metadata describing the headers and provides the metadata to the downstream functional components, such as the header lookup **106**. For example, the parser **104** updates the metadata item to include data identifying the header types of the individual headers in the header layer, as well as the location of the header layer and/or the individual headers included in the header layer. The location of a header layer or an individual header may be defined by a range pair that includes an offset value identifying a position (e.g., bit position) at which the header layer or header begins in the data unit (e.g., header byte string), as well as a width value identifying a length (e.g., number of bits) of the header layer or header. The

metadata item may include an array of header types and range pairs, which the parser **104** updates to include the determined header types in the header layer and their corresponding range pairs.

(30) The parser **104** provides the metadata item, including the array of header types and range pairs, to the lookup result processor **108**. The parser **104** may also provide the header lookup **106** with the header byte string, such as by providing the header lookup **106** with the header byte string alone, header byte string along with some portion of the data unit following the header byte string, or the entirety of the data unit. Providing the downstream functional components (e.g., header lookup **106**) with metadata describing the headers, rather than data extracted from the headers, allows each functional component to determine which header fields to use when performing any subsequent operations. This approach provides flexibility while also yielding a compact and efficient implementation.

(31) FIG. 2 is a block diagram showing a processing loop **200** performed by the parser **104**, according to some example embodiments. As explained earlier, headers do not include encoded values that identify the header's type, but rather a previous (e.g., outer) header includes data indicating the header type of the next sequential header in a data unit. As the outermost (e.g., first) header does not have a previous header from which to access a next header type, the parser **104** uses data identifying the port via which the data unit was received (e.g., unique port identifier) to infer the header type of the outermost header. For example, the parser **104** accesses the port identifier from the metadata item received from the receive port manager **102**.

(32) As shown in FIG. 2, the port identifier is provided as input into an initial header types table **202**. The initial header types table **202** includes a listing of port identifiers mapped to header type values. The port identifier can therefore be used as an index into the initial header types table **202** to cause output of the corresponding header type value. The returned header type value is inferred to identify the header type of the outermost header.

(33) A header type value can be non-enumerated (e.g., there are no predefined values for various header types, completely enumerated (e.g., all supported headers have pre-defined header type values), or partially enumerated (e.g., a value denoting “unknown” or “user data” may be pre-defined).

(34) The header type value may be derived from either the port identifier when processing the outermost header or a next header field of a preceding header when processing any subsequent headers. Accordingly, a multiplexer **204** is used to select the proper source of the header type value that is provided to downstream components of the processing loop **200**. As shown, the multiplexer **204** receives input from both the initial header types table **202** and the next header types table **212**, as well as a first header indicator. The first header indicator acts as a control to the multiplexer **204** to cause the multiplexer to select either the output of the initial header types table **202** (e.g., when processing the outermost header) or the output of the next-header types table **212** (e.g., when processing any subsequent header).

(35) The header type value outputted by multiplexer **204** is provided as an input into the field extraction parameters table **206**. The field extraction parameters table **206** maintains a listing of range pairs (e.g., offset value and width value) corresponding to various header type values. Accordingly, the header type value is used as an index or search argument into the field extraction parameters table **206** to cause output of the offset value and width value corresponding to the next-header type value.

(36) The offset value and width value are provided as input into the field extractor **208**. The field extractor **208** uses the offset value and width value to extract the next-header field from the header byte string **210**. For example, the offset value and width value are used to control the logic of the field extractor **208** to cause the field extractor **208** to extract the specified next-header field from the header byte string **210**. The next-header field identifies the header type of the subsequent header in the data unit, if any.

(37) The next-header field extracted by the field extractor **208** is concatenated with the header type

value outputted by the multiplexer **204** and the resulting concatenated data is provided as input into the next-header types table **212**. The next-header types table **212** includes a listing of header types corresponding to various inputs of concatenated header type values and next-header field data. The concatenated data is used to form an index or a search argument into the next-header types table **212** to cause output of a header type corresponding to the input data. The concatenation of the header type value outputted by the multiplexer **204** and next-header field value extracted from the header byte string **210** is used to identify the proper encoding used by header. For example, both Ethernet and VXLAN use 16-bit next-header fields, however the encodings used by each are completely different. Accordingly, the header type identifies the proper encoding, which allows for proper interpretation of the data included in the next-header field.

(38) The header type value that is provided as output by the next-header types table **212** can be used as both an output of the processing loop **200** and as an input to the multiplexer to cause parsing of the subsequent header in the header byte string. For example, the first header indicator can control the multiplexer to select the header type value provided by the next-header types table **212** as input into the field extraction parameters table **206**, to cause reiteration of the functionality of the processing loop **200** based on subsequent header, such as determining the range pair of the subsequent header, extracting and next-header field data from the subsequent header, and determining a header type based on the next-header field data extracted from the subsequent header.

(39) Returning to the discussion of FIG. **1**, the parser **104** and the header lookup **106** work together to process the data unit header-by-header. For example, the parser **104** is followed in the data processing pipeline by the header lookup **106**, and the header lookup **106** is followed by the parser **104** along the recirculation path. Each pass through this header processing loop **116** performed by the parser **104** and header lookup **106** processes or “consumes” a single layer of headers. In each pass through the header processing loop **116**, the parser **104** provides the header lookup **106** with an updated metadata item that includes data describing a subsequent layer of headers included in the data unit, such as the header types and range pair for each header in the layer of headers. In turn, the header lookup **106** further processes the data unit based on the subsequent layer of headers.

(40) Once the header lookup **106** is done processing a layer of headers, it may return the data unit to the parser **104** via the recirculation path to cause the parser **104** to parse the next layer of headers. The data unit returned to the parser **104** may be represented by just the metadata item, the metadata item and the header data byte string, or by the metadata item and the entire data unit. Returning the data unit to the parser **104** is typically performed when the header lookup **106** determines that a header represents a tunnel that is being terminated at the current network node. When terminating a tunnel, the payload of a data unit is processed as if it was a newly received data unit. Accordingly, the subsequent layer of headers is parsed and processed in the same general manner as the current layer. Thus, the relationship that the parser **104** and the header lookup **106** share via the recirculation path provides the protocol independent data unit forwarding system **100** with the flexibility to efficiently terminate tunnels, as well as to perform other operations, such as update processing contexts, determine destinations, as described below.

(41) A single, flexible search argument assembly and lookup function supports the various operations performed by the header lookup **106**, such as refining the data unit's context (e.g., updating the metadata item), managing the contents of forwarding tables, testing the data unit against access control lists, determining the data unit's destination and/or forwarding requirements, hashing flow-identifying entropy for load balancing purposes, and the like. This flexibility is achieved by using the non-enumerated layer type values generated by the parser **104** to generate search arguments that are used to perform lookup operations in tables that hold keys that are compatible with the assembled search arguments. For example, the header lookup **106** uses data included in the metadata item to retrieve sets of parameters that are used to configure bit-field

extractors. Examples of data included in the metadata item include layer type values, offset values, port identifier for a physical receive port, port identifier for a logical receive port, layer 2 forwarding domain, layer 3 forwarding domain, header type, quality of service code points (e.g., priority), address value type (e.g., unicast, multicast, etc.), and the like. The configured bit-field extractors are then used to concatenate arbitrary strings of header data into the search arguments, which are used to perform the lookup operations.

(42) Different types of lookup operations (e.g., search algorithms) may be used for different purposes. For example, an indexed read lookup operation may be used to map a data unit's port identifier to the destination identifier. As another example, an exact-match search lookup operation may be used to bridge Ethernet data units (e.g., Ethernet packets) based on their headers. As another example, a longest-prefix match lookup operation may be used to route data based on their headers (e.g., IPv4 or IPv6). Generation of search arguments and performing lookup operations in this manner enables the header lookup **106** to operate freely of fixed relationships between a data unit's forwarding header type, forwarding method, or the lookup algorithms used to make forwarding or other decisions based on header contents.

(43) Each lookup operation, whether successful or not, returns some data. For example, successful lookups (e.g., those where a matching key is found) return data that is associated with the matching key value. Unsuccessful lookups return default data that is associated with the table in which the lookup operation was performed. In either case, the data returned as a result of a lookup operation can be used to determine a subsequent action performed by the header lookup **106**.

(44) At the completion of each individual lookup operation, the header lookup **106** performs a three-way branch decision to either perform another lookup operation, recirculate the data unit back to the parser **104** to parse the subsequent layer of headers, or terminate the header processing loop **116** and provide an output to the lookup result processor **108**. The header lookup **106** may determine the appropriate subsequent action based on the output of the previous lookup. For example, the output of a lookup operation may include data identifying a next lookup operation to be performed, such as a pointer to the next lookup operation. This may be used to update the metadata item using lookup operations on the fields of one or more headers.

(45) As another example, the output of a lookup operation may include data indicating that a subsequent layer of headers should be parsed, necessitating recirculation of the data unit back to the parser **104**. For example, this capability may be used to terminate a tunnel.

(46) As another example, the output of a lookup operation may include data instructing the header lookup **106** to terminate the header processing loop **116** and by providing output to the lookup result processor **108**. The output provided to the lookup result processor may include the metadata item, metadata item and the header byte string, or the metadata item and entire data unit or a portion of the data unit that is greater than the header byte string.

(47) Prior to providing output to the lookup result processor **108** the header lookup **106** determines a first forwarding instruction to be performed to forward the data unit to its destination. The header lookup **106** determines the first forwarding instruction during performance of the lookup operations described above. The header lookup **106** provides the lookup result processor **108** with data identifying the first forwarding instruction, such as by updating the metadata item to include a value that points to the first instruction. This may be performed using an action instruction that is returned as a result of a destination-determining lookup operation. However, an action associated with a tunnel termination lookup operation or an access control list operation may also be used to set value that points to the first instruction. Similarly, other methods for setting the value that points to the first instruction may also be used.

(48) The lookup result processor **108** uses the value that points to the first instruction to determine a finalized forwarding decision. The forwarding decision may incorporate such behaviors as load balancing, multicast replication, and the like. A number of other functions may be performed by the lookup result processor **108**, such as output queue selection. Alternatively, the forwarding decision

may be limited to identifying the transmit port to be used to forward the data unit to its intended destination. Typically, a single transmit port may be served by multiple output queues, each with its own relative priority, bandwidth limits, and other quality of service-related characteristics. After identifying the transmit port for the data unit, the lookup result processor **108** may use data included in the metadata item (e.g., the context of the data unit) and values conveyed by its priority related field to select one of the output queues that service the transmit port. After the output queue is selected, the data unit is ready to be enqueued to the output queue for eventual transmission of the data unit.

(49) The data buffer **110** provides temporary storage for a data unit, the queuing of a data unit for retrieval and transmission, and mechanisms for arbitrating among the output queues and implementing specific byte-rate limiting and byte-rate guarantees. In embodiments in which the header byte string is delivered to the data buffer **110** separately from the body of the data unit, the data buffer may also support the merging of the header data byte string with the body byte string to form a complete data unit. Although the data buffer **110** is positioned prior to the header rewrite **112**, this is only one embodiment and is not meant to be limiting. In other embodiments, the data buffer **110** may be positioned after the header rewrite **112** in the data processing pipeline.

(50) The header rewrite **112** is responsible for updating the header byte string for a data unit to conform with the requirements of the forwarding method being employed to transmit the data unit and/or the network path that the data unit will follow (e.g., VLANs, tunnels, etc.). Specifically, the header rewrite **112** may strip the encapsulating headers that are associated with terminated tunnels, update the forwarding header, and append new encapsulating headers.

(51) In a simple example, such as VLAN-unaware Ethernet bridging, zero bytes are stripped from the data unit, no changes are made to the forwarding header, and no new encapsulating headers are appended to the data unit. In a more complex example, the any combination of the three actions listed above may be performed by the header rewrite **112**.

(52) In some embodiments, the header rewrite **112** uses data included in the metadata item to remove the encapsulating headers that are associated with terminated tunnels. For example, the header lookup **106** may have determined which header is serving as the forwarding header and, the header rewrite **112** may use the offset value for that header to remove the encapsulating headers, such as by removing the headers that located before the offset value.

(53) The header rewrite **112** uses algorithmic modifications to update a forwarding header. For example, the header rewrite **112** performs algorithmic modifications to one or more header fields, such as by decrementing a time-to-live value or by checksum computation. As another example, the header rewrite **112** may translate or map a header field value to another value, such as VLAN ID translation or network address translation.

(54) The header rewrite **112** can use templates to generate new encapsulating headers for the data unit. For example, the header rewrite **112** customizes a template with tunnel-specific values such as source and destination addresses, a time-to-live value, and quality of service code points. The operations performed by the header rewrite **112** may either occur within a buffer that holds the entire data unit or within a separate buffer (e.g., data buffer **110**) that holds the header byte string (e.g., containing, at least, the forwarding header). If a separate buffer is used, the updated headers are appended to the remainder of the data unit to form a complete data unit.

(55) The transmit port manager **114** facilitates transmit operations. Although the transmit port manager **114** and the receive port manager **102** are shown as separate functional components, this is just one example and is not meant to be limiting. In some embodiments, the functionality of the transmit port manager **114** and the receive port manager **102** are embodied within a single functional component. The transmit port manager **114** embodies various MACs and operates to transmit data units to other devices in a network via an assortment of external ports. The external ports may include Ethernet, as well as a variety of specialty ports such as CSI, CAN, LIN, and the like. Each port is identified by a unique port identifier.

(56) The transmit port manager **114** may perform a variety of data processing operations to prepare a data unit for transmission to another device. For example, the transmit port manager **114** may perform any media specific adaptations on the received data unit as well as demultiplex the data unit from the data processing pipeline.

(57) FIG. **3** is a flowchart showing a method **300** for receiving and initial processing of a data unit for protocol independent data unit forwarding, according to some example embodiments. The method **300** may be embodied in computer readable instructions for execution by one or more processors such that the operations of the method **300** may be performed in part or in whole by the receive port manager **102**; accordingly, the method **300** is described below by way of example with reference to the receive port manager **102**. However, it shall be appreciated that at least some of the operations of the method **300** may be deployed on various other hardware and/or software configurations and the method **300** is not intended to be limited to the receive port manager **102**.

(58) At operation **302**, the receive port manager **102** receives a data unit. The receive port manager **102** embodies various MACs and operates to receive data units transmitted in a network via an assortment of external ports. The external ports may include Ethernet, as well as a variety of specialty ports such as CSI, CAN, LIN, and the like. Each port is identified by a unique port identifier.

(59) After receiving a data unit, the receive port manager **102** may perform a variety of data processing operations to generate an output that can be provided to the other functional components of the protocol independent data unit forwarding system **100**. For example, at operation **304**, the receive port manager **102** performs an adaptation, such as any media specific adaptations on the received data unit. Similarly, at operation **306**, the receive port manager **102** multiplexes the data unit onto a common multiplexed path so that the data unit, or any portion, thereof may be transmitted to the other downstream functional components of the protocol independent data unit forwarding system **100**.

(60) At operation **308**, the receive port manager **102** identifies the header byte string from the data unit. The header byte string includes a subset of the ordered set of individual values from the data unit. For example, the header byte string may include a predetermined number of the first ordered values (e.g., first *n* bytes) from the data unit. The predetermined number may be any number that is less than or equal to the total number of values in the data unit, however the predetermined number used may be selected to encompass all relevant header data included in the data unit, while remaining small enough that the size of the header byte string does not provide undue burden during subsequent processing of the header byte string.

(61) At operation **310**, the receive port manager **102** generates a metadata item for the data unit. The generated metadata item for the data unit can be used to store data that describes the context of the data unit. For example, the metadata item includes a series of values that describe a context of the data unit that can be used by the various functional components of the protocol independent data unit forwarding system **100** to process the data unit. A functional component may use the values in the metadata item to determine which operation the functional component is to perform to further process the data unit. As each individual operation is performed, the metadata item may be updated to further define the context of the data unit. For example, the metadata item may be updated based on the resulting output of an operation by adding new values to the metadata item, modifying existing values in the metadata item, and the like.

(62) The receive port manager **102** configures some initial data values in the metadata item. For example, the receive port manager **102** may include data values in the metadata item to identify the port via which the data unit was received (e.g., unique port identifier), a data length of the data unit (e.g., byte count), a time stamp of the time at which the data unit was received, and/or data identifying any errors that may have been detected during reception of the data unit.

(63) At operation **312**, the receive port manager **102** forwards the data unit and the metadata item to downstream operational components of the protocol independent data unit forwarding system **100**.

For example, the receive port manager **102** may provide the parser **104** with only the header byte string and the metadata item. In this type of embodiment, the receive port manager **102** may copy the portion of the data unit following the header byte string to the data buffer **110** to bypass the parser **104**, header lookup **106**, and lookup result processor **108**. As a result, the portion of the data unit following the header byte string is not transmitted through parser **104**, header lookup **106**, and lookup result processor **108**, which reduces overall resource usage associated with processing the data unit.

(64) In other embodiments, however, the receive port manager **102** may transmit the metadata item and the entire data unit or a portion of the data unit that is greater than the header byte string to the parser **104**. Further, in some embodiments, the receive port manager **102** may transmit the entire data unit or the portion of the data unit following the header byte string to the data buffer **110** along with some portion of the header byte string to the data buffer **110**. Alternatively, the receive port manager **102** may not transmit any data to the data buffer **110**, such as in embodiments in which the receive port manager **102** provides the entire data unit to the parser **104**. In this type of embodiment, the protocol independent data unit forwarding system **100** may not include the bypass path to the data buffer **110**.

(65) FIG. **4** is a flowchart showing a method **400** for parsing a data unit for protocol independent data unit forwarding, according to some example embodiments. The method **400** may be embodied in computer readable instructions for execution by one or more processors such that the operations of the method **400** may be performed in part or in whole by receive the parser **104**; accordingly, the method **400** is described below by way of example with reference to the parser **104**. However, it shall be appreciated that at least some of the operations of the method **400** may be deployed on various other hardware and/or software configurations and the method **400** is not intended to be limited to the parser **104**.

(66) At operation **402**, the parser **104** determines a header type of a current header included in the data unit. The parser **104** may determine the header type in different ways based on whether the header is the outermost header or any subsequent header in a data unit. In general, headers do not include encoded values that identify the header's type. Rather, a previous (e.g., outer) header indicates the header type of the next sequential header in the data unit. Accordingly, for any header that is subsequent to the outermost header, the parser **104** determines the header type based on the next header field of the preceding header.

(67) As the outermost (e.g., first) header does not have a previous header from which to access a next header type, the parser **104** uses the metadata item received from the receive port manager **102** to identify the header type of the outermost (e.g., first) header. Specifically, the parser **104** identifies the header type of the outermost header based on the data identifying the port via which the data unit was received (e.g., unique port identifier). For example, each port or port identifier may be mapped to a header type. Accordingly, the parser **104** determines the header type mapped to the unique port identifier received in the metadata items, which is inferred to be the header type of the outermost header. The mapping between the ports and the header types of the outermost header allows the parser **104** to operate with any of a variety of network types.

(68) At operation **404**, the parser **104** determines a range pair identifying a location of the current header in the data unit. The range pair includes an offset value identifying a position (e.g., bit position) at which the header layer or header begins in the data unit (e.g., header byte string), as well as a width value identifying a length (e.g., number of bits) of the header layer or header. The parser determines the range pair based on the header type of the header. For example, the header type (e.g., header type value) is used as an index or search argument into a field extraction parameters table **206** to cause output of the offset value and width value corresponding to the header type value. The field extraction parameters table **206** maintains a listing of range pairs (e.g., offset value and width value) corresponding to various header type values.

(69) At operation **406**, the parser **104** extracts a next-header field from the current header based on

the range pair. For example, the parser **104** uses the offset value and width value to control the logic of a field extractor **208** to cause the field extractor **208** to extract the specified next-header field from the header byte string **210**. The next-header field identifies the header type of the subsequent header in the data unit, if any.

(70) At operation **408**, the parser **104** determines a next header type based on the header type of the current header and the next-header field. For example, the parser **104** concatenates the next-header field with the header type value and the resulting concatenated data is provided as input into a next header types table **212**. The next header types table **212** includes a listing of header types corresponding to various inputs of concatenated header type values and next-header field data. The concatenated data is used to form an index or a search argument into the next header types table **212** to cause output of a header type corresponding to the input data.

(71) At operation **410**, the parser **104** determines whether to process the next header. For example, the parser **104** may determine whether the current header and the next header are should be included in a single layer of headers that should be processed together by the downstream components of the protocol independent data unit forwarding system **100**. A layer of headers is generally a collection of one or more headers that are closely related in a manner such that information from each of the headers is used to analyze or understand the data included in the collection headers, although the definition of a header layer may be configured as desired such that a header layer does not necessarily have to include headers that are related in this manner. For example, a header layer may include headers that are not related as described above and/or headers that are related as described above may be split into separate header layers.

(72) At operation **412**, the parser **104** provides output data to the header lookup **106**. For example, the parser **104** may provide the header lookup **106** with the metadata item associated with the data unit. The metadata item may include an array of header types and range pairs includes data describing the layers of headers included in the data unit that have been identified by the parser **104**. The parser **104** may also provide the header lookup **106** with the header byte string, such as by providing the header lookup **106** with the header byte string alone, header byte string along with some portion of the data unit following the header byte string, or the entirety of the data unit.

(73) FIG. 5 is a flowchart showing a method **500** of performing a header lookup for protocol independent data unit forwarding, according to some example embodiments. The method **500** may be embodied in computer readable instructions for execution by one or more processors such that the operations of the method **500** may be performed in part or in whole by receive the header lookup **106**; accordingly, the method **500** is described below by way of example with reference to the header lookup **106**. However, it shall be appreciated that at least some of the operations of the method **500** may be deployed on various other hardware and/or software configurations and the method **500** is not intended to be limited to the header lookup **106**.

(74) At operation **502**, the header lookup **106** receives a metadata item associated with a data unit. The metadata item is received from the parser **104** and includes data describing a context of the data unit. For example, the metadata item may include data describing layers of headers included in the data unit (e.g., header types and range pair), port identifier for a physical receive port, port identifier for a logical receive port, layer **2** forwarding domain, layer **3** forwarding domain, header type, quality of service code points (e.g., priority), address value type (e.g, unicast, multicast, etc.), and the like.

(75) At operation **504**, the header lookup **106** generates a search argument based on the metadata item. For example, the header lookup **106** uses data included in the metadata item to retrieve a set of parameters that are used to configure bit-field extractors. The parameters may be used to identify a table into which the lookup operation is be performed and select the lookup operation type (e.g., search algorithm) to be used. These parameters may also specify a list of metadata identifiers and a series of offset/width pairs that are used to extract arbitrary strings of bits from the header byte string. The selected metadata and the extracted bit strings are then concatenated to form a search

argument.

(76) At operation **506**, the header lookup **106** executes a lookup based on the search argument. For example, the search argument is used to perform a lookup operation of the selected lookup operation type in the identified table. Each lookup operation, whether successful or not, returns some data. For example, successful lookups (e.g., those where a matching key is found) return data that is associated with the matching key value. Unsuccessful lookups return default data that is associated with the table in which the lookup operation was performed and/or is associated with the lookup instruction used to initiate the lookup operation.

(77) At operation **508**, the header lookup **106** determines a subsequent action based on the lookup output. For example, the data returned as a result of a lookup operation can identify or be used to determine the subsequent action performed by the header lookup **106**. As shown, the header lookup **106** performs a three-way branch decision to either perform another lookup operation, recirculate the data unit back to the parser **104** to parse the subsequent layer of headers, or terminate the header processing loop **116** and provide an output to the lookup result processor **108**. The header lookup **106** may determine the appropriate subsequent action based on the output of the previous lookup. For example, the output of a lookup operation may include data identifying the subsequent action.

(78) If the header lookup up determines to recirculate the data unit back to the parser **104**, the method **500** proceeds to operation **510**, where the header lookup **106** recirculates the data unit to the parser **104**. For example, this capability may be used to terminate a tunnel.

(79) Alternatively, if the header lookup up determines to perform another lookup operation, the method **500** proceeds to operation **512**, where the header lookup **106** performs another lookup operation. This may be used to update the metadata item using lookup operations on the fields of one or more headers.

(80) Alternatively, if the header lookup up determines to terminate the header processing loop **116** and provide an output to the lookup result processor **108**, the method **500** proceeds to operation **514**, where the header lookup **106** provides output to the lookup result processor **108**. The output provided to the lookup result processor **108** may include the metadata item, metadata item and the header byte string, or the metadata item and entire data unit or a portion of the data unit that is greater than the header byte string.

(81) Prior to providing output to the lookup result processor **108** the header lookup **106** determines a first forwarding instruction to be performed to forward the data unit to its destination. The header lookup **106** determines the first forwarding instruction during performance of the lookup operations described above. The header lookup **106** provides the lookup result processor **108** with data identifying the first forwarding instruction, such as by updating the metadata item to include a value that points to the first instruction. This may be performed using an action instruction that is returned as a result of a destination-determining lookup operation. However, an action associated with a tunnel termination lookup operation or an access control list operation may also be used to set a value that points to the first instruction. Similarly, other methods for setting the value that points to the first instruction may also be used.

(82) FIG. **6** is a flowchart showing a method **600** protocol independent data unit forwarding, according to some example embodiments. The method **600** may be embodied in computer readable instructions for execution by one or more processors such that the operations of the method **600** may be performed in part or in whole by a protocol independent data unit forwarding system **100**; accordingly, the method **600** is described below by way of example with reference to the protocol independent data unit forwarding system **100**. However, it shall be appreciated that at least some of the operations of the method **600** may be deployed on various other hardware and/or software configurations and the method **600** is not intended to be limited to the protocol independent data unit forwarding system **100**.

(83) At operation **602**, the receive port manager **102** receives a data unit. The data unit is received

via an input port of the protocol independent packet forwarding system **100**. The data unit comprises a header byte string that is a subset of the data unit.

(84) At operation **604**, the parser **104** parses the data unit based on the input port. For example, the parser determines a header type based on the input port and then determines a range pair defining a set of bits from the header byte string that represent an outermost header in the header byte string. In some embodiments, the parser **104** may also determine another range pair defining a set of bits from the header byte string that represent a subsequent header in the header byte string.

(85) At operation **606**, the parser **104** updates a metadata item associated with the data unit based on a parsing output. For example, the parser updates the metadata item to include the range pairs and header types.

(86) At operation **608**, the header lookup **106** determines a packet forwarding instruction based on the metadata item. For example, the header lookup **106** may generate a search argument based on the metadata item, and execute a lookup based on the search argument. The header lookup **106** may then determine a subsequent action based on the lookup output.

(87) In some embodiments, the subsequent action is performing a subsequent parsing of the data unit. In this case, the header lookup **106** may instruct the parser **104** to parse the data unit based on a subsequent header type determined based on the lookup output. As a result, the parser **104** updates the metadata item associated with the data unit based on a subsequent parsing output. The header lookup **106** may then generate a subsequent search argument based on the updated metadata item and execute a subsequent lookup based on the subsequent search argument. The header lookup may determine another subsequent action based on the lookup output of the subsequent lookup.

(88) In some embodiments, the subsequent action is to perform another lookup. In this case, the header lookup **106** may update the metadata item associated with the data unit based on the lookup output and generate a subsequent search argument based on the updated metadata item. The header lookup **106** may then execute a subsequent lookup based on the subsequent search argument and determine another subsequent action based on the lookup output of the subsequent lookup.

(89) In some embodiments, the subsequent action is to forward the data unit to the first destination. For example, the lookup output may identify the packet forwarding instruction, which the header lookup **106** can provide to the lookup result processor **108**.

(90) At operation **610**, the protocol independent data unit forwarding system **100** forwards the data unit to its destination based on the forwarding instruction. For example, the data unit may be forwarded according to the method **700** shown in FIG. 7.

(91) FIG. 7 is a flowchart showing a method **700** of forwarding a data unit to its destination based on a forwarding instruction, according to some example embodiments. The method **700** may be embodied in computer readable instructions for execution by one or more processors such that the operations of the method **700** may be performed in part or in whole by a protocol independent data unit forwarding system **100**; accordingly, the method **700** is described below by way of example with reference to the protocol independent data unit forwarding system **100**. However, it shall be appreciated that at least some of the operations of the method **700** may be deployed on various other hardware and/or software configurations and the method **700** is not intended to be limited to the protocol independent data unit forwarding system **100**.

(92) At operation **702**, the lookup result processor **108** determines a transmit port for a data unit based on a packet forwarding instruction. Typically, a single transmit port may be served by multiple output queues, each with its own relative priority, bandwidth limits, and other quality of service-related characteristics. Accordingly, after identifying the transmit port for the data unit, at operation **704**, the lookup result processor **108** selects an output queue servicing the transmit port based on a metadata item associated with the data unit. For example, the lookup result processor **108** may use data included in the metadata item (e.g., the context of the data unit) and values conveyed by its priority related field to select one of the output queues that service the transmit port.

(93) At operation **706**, the header rewrite **112** generates an updated header byte string for the data unit. The header rewrite **112** is responsible for updating the header byte string for a data unit to conform with the requirements of the forwarding method being employed to transmit the data unit and/or the network path that the data unit will follow (e.g., VLANs, network tunnels, etc.).

Specifically, the header rewrite **112** may strip the encapsulating headers that are associated with terminated tunnels, update the forwarding header; and append new encapsulating headers.

(94) In a simple example, such as VLAN-unaware Ethernet bridging, zero bytes are stripped from the data unit, no changes are made to the forwarding header, and no new encapsulating headers are appended to the data unit. In a more complex example, the any combination of the three actions listed above may be performed by the header rewrite **112**.

(95) In some embodiments, the header rewrite **112** uses data included in the metadata item to remove the encapsulating headers that are associated with terminated tunnels. For example, the header lookup **106** may have determined which header is serving as the forwarding header and, the header rewrite **112** may use the offset value for that header to remove the encapsulating headers, such as by removing the headers located before the offset value.

(96) The header rewrite **112** uses algorithmic modifications to update a forwarding header. For example, the header rewrite **112** performs algorithmic modifications to one or more header fields, such as by decrementing a time-to-live value or by checksum computation. As another example, the header rewrite **112** may translate or map a header field value to another value, such as a VLAN ID translation or network address translation.

(97) The header rewrite **112** can use templates to generate new encapsulating headers for the data unit. For example, the header rewrite **112** customizes a template with tunnel-specific values such as source and destination addresses, a time-to-live value, and quality of service code points. The operations performed by the header rewrite **112** may either occur within a buffer that holds the entire data unit or within a separate buffer (e.g., data buffer **110**) that holds the header byte string (e.g., containing, at least, the forwarding header). If a separate buffer is used, the updated headers is appended to the remainder of the data unit to form a complete packet.

(98) At operation **708**, the header rewrite **112** adds the data unit to the output queue servicing the transmit port.

(99) Software Architecture

(100) FIG. **8** is a block diagram illustrating an example software architecture **806**, which may be used in conjunction with various hardware architectures herein described. FIG. **8** is a non-limiting example of a software architecture **806** and it will be appreciated that many other architectures may be implemented to facilitate the functionality described herein. The software architecture **806** may execute on hardware such as machine **900** of FIG. **9** that includes, among other things, processors **904**, memory **914**, and (input/output) I/O components **918**. A representative hardware layer **852** is illustrated and can represent, for example, the machine **900** of FIG. **9**. The representative hardware layer **852** includes a processing unit **854** having associated executable instructions **804**. Executable instructions **804** represent the executable instructions of the software architecture **806**, including implementation of the methods, components, and so forth described herein. The hardware layer **852** also includes memory and/or storage modules **856**, which also have executable instructions **804**. The hardware layer **852** may also comprise other hardware **858**.

(101) In the example architecture of FIG. **8**, the software architecture **806** may be conceptualized as a stack of layers where each layer provides particular functionality, such as the Open Systems Interconnection model (OSI model). For example, the software architecture **806** may include layers such as an operating system **802**, libraries **820**, frameworks/middleware **818**, applications **816**, and a presentation layer **814**. Operationally, the applications **816** and/or other components within the layers may invoke application programming interface (API) calls **808** through the software stack and receive a response such as messages **812** in response to the API calls **808**. The layers illustrated are representative in nature and not all software architectures have all layers. For example, some

mobile or special purpose operating systems may not provide a frameworks/middleware **818**, while others may provide such a layer. Other software architectures may include additional or different layers.

(102) The operating system **802** may manage hardware resources and provide common services. The operating system **802** may include, for example, a kernel **822**, services **824**, and drivers **826**. The kernel **822** may act as an abstraction layer between the hardware and the other software layers. For example, the kernel **822** may be responsible for memory management, processor management (e.g., scheduling), component management, networking, security settings, and so on. The services **824** may provide other common services for the other software layers. The drivers **826** are responsible for controlling or interfacing with the underlying hardware. For instance, the drivers **826** include display drivers, camera drivers, Bluetooth® drivers, flash memory drivers, serial communication drivers (e.g., Universal Serial Bus (USB) drivers), Wi-Fi® drivers, audio drivers, power management drivers, and so forth, depending on the hardware configuration.

(103) The libraries **820** provide a common infrastructure that is used by the applications **816** and/or other components and/or layers. The libraries **820** provide functionality that allows other software components to perform tasks in an easier fashion than to interface directly with the underlying operating system **802** functionality (e.g., kernel **822**, services **824**, and/or drivers **826**). The libraries **820** may include system libraries **844** (e.g., C standard library) that may provide functions such as memory allocation functions, string manipulation functions, mathematical functions, and the like. In addition, the libraries **820** may include API libraries **846** such as media libraries (e.g., libraries to support presentation and manipulation of various media format such as MPEG4, H.264, MP3, AAC, AMR, JPG, PNG), graphics libraries (e.g., an OpenGL framework that may be used to render 2D and 3D in a graphic content on a display), database libraries (e.g., SQLite that may provide various relational database functions), web libraries (e.g., WebKit that may provide web browsing functionality), and the like. The libraries **820** may also include a wide variety of other libraries **848** to provide many other APIs to the applications **816** and other software components/modules.

(104) The frameworks/middleware **818** (also sometimes referred to as middleware) provide a higher-level common infrastructure that may be used by the applications **816** and/or other software components/modules. For example, the frameworks/middleware **818** may provide various graphical user interface (GUI) functions, high-level resource management, high-level location services, and so forth. The frameworks/middleware **818** may provide a broad spectrum of other APIs that may be used by the applications **816** and/or other software components/modules, some of which may be specific to a particular operating system **802** or platform.

(105) The applications **816** include built-in applications **838** and/or third-party applications **840**. Examples of representative built-in applications **838** may include, but are not limited to, a contacts application, a browser application, a book reader application, a location application, a media application, a messaging application, and/or a game application. Third-party applications **840** may include an application developed using the ANDROID™ or IOS™ software development kit (SDK) by an entity other than the vendor of the particular platform, and may be mobile software running on a mobile operating system such as IOS™, ANDROID™, WINDOWS® Phone, or other mobile operating systems. The third-party applications **840** may invoke the API calls **808** provided by the mobile operating system (such as operating system **802**) to facilitate functionality described herein.

(106) The applications **816** may use built in operating system functions (e.g., kernel **822**, services **824**, and/or drivers **826**), libraries **820**, and frameworks/middleware **818** to create UIs to interact with users of the system. Alternatively, or additionally, in some systems, interactions with a user may occur through a presentation layer, such as presentation layer **814**. In these systems, the application/component “logic” can be separated from the aspects of the application/component that interact with a user.

(107) FIG. 9 is a block diagram illustrating components of a machine **900**, according to some

example embodiments, able to read instructions **804** from a machine-readable medium (e.g., a machine-readable storage medium) and perform any one or more of the methodologies discussed herein. Specifically, FIG. **9** shows a diagrammatic representation of the machine **900** in the example form of a computer system, within which instructions **910** (e.g., software, a program, an application, an applet, an app, or other executable code) for causing the machine **900** to perform any one or more of the methodologies discussed herein may be executed. As such, the instructions **910** may be used to implement modules or components described herein. The instructions **910** transform the general, non-programmed machine **900** into a particular machine **900** programmed to carry out the described and illustrated functions in the manner described. In alternative embodiments, the machine **900** operates as a standalone device or may be coupled (e.g., networked) to other machines. In a networked deployment, the machine **900** may operate in the capacity of a server machine or a client machine in a server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine **900** may comprise, but not be limited to, a server computer, a client computer, a PC, a tablet computer, a laptop computer, a netbook, a set-top box (STB), a personal digital assistant (PDA), an entertainment media system, a cellular telephone, a smart phone, a mobile device, a wearable device (e.g., a smart watch), a smart home device (e.g., a smart appliance), other smart devices, a web appliance, a network router, a network switch, a network bridge, or any machine **900** capable of executing the instructions **910**, sequentially or otherwise, that specify actions to be taken by machine **900**. Further, while only a single machine **900** is illustrated, the term “machine” shall also be taken to include a collection of machines that individually or jointly execute the instructions **910** to perform any one or more of the methodologies discussed herein.

(108) The machine **900** may include processors **904**, memory/storage **906**, and I/O components **918**, which may be configured to communicate with each other such as via a bus **902**. The memory/storage **906** may include a memory **914**, such as a main memory, or other memory storage, and a storage unit **916**, both accessible to the processors **904** such as via the bus **902**. The storage unit **916** and memory **914** store the instructions **910** embodying any one or more of the methodologies or functions described herein. The instructions **910** may also reside, completely or partially, within the memory **914**, within the storage unit **916**, within at least one of the processors **904** (e.g., within the processor's cache memory), or any suitable combination thereof, during execution thereof by the machine **900**. Accordingly, the memory **914**, the storage unit **916**, and the memory of processors **904** are examples of machine-readable media.

(109) The I/O components **918** may include a wide variety of components to receive input, provide output, produce output, transmit information, exchange information, capture measurements, and so on. The specific I/O components **918** that are included in a particular machine **900** will depend on the type of machine. For example, portable machines such as mobile phones will likely include a touch input device or other such input mechanisms, while a headless server machine will likely not include such a touch input device. It will be appreciated that the I/O components **918** may include many other components that are not shown in FIG. **9**. The I/O components **918** are grouped according to functionality merely for simplifying the following discussion and the grouping is in no way limiting. In various example embodiments, the I/O components **918** may include output components **926** and input components **928**. The output components **926** may include visual components (e.g., a display such as a plasma display panel (PDP), a light emitting diode (LED) display, a liquid crystal display (LCD), a projector, or a cathode ray tube (CRT)), acoustic components (e.g., speakers), haptic components (e.g., a vibratory motor, resistance mechanisms), other signal generators, and so forth. The input components **928** may include alphanumeric input components (e.g., a keyboard, a touch screen configured to receive alphanumeric input, a photo-optical keyboard, or other alphanumeric input components), point based input components (e.g., a mouse, a touchpad, a trackball, a joystick, a motion sensor, or other pointing instrument), tactile input components (e.g., a physical button, a touch screen that provides location and/or force of

touches or touch gestures, or other tactile input components), audio input components (e.g., a microphone), and the like.

(110) In further example embodiments, the I/O components **918** may include biometric components **930**, motion components **934**, environmental components **936**, or position components **938** among a wide array of other components. For example, the biometric components **930** may include components to detect expressions (e.g., hand expressions, facial expressions, vocal expressions, body gestures, or eye tracking), measure biosignals (e.g., blood pressure, heart rate, body temperature, perspiration, or brain waves), identify a person (e.g., voice identification, retinal identification, facial identification, fingerprint identification, or electroencephalogram based identification), and the like. The motion components **934** may include acceleration sensor components (e.g., accelerometer), gravitation sensor components, rotation sensor components (e.g., gyroscope), and so forth. The environmental components **936** may include, for example, illumination sensor components (e.g., photometer), temperature sensor components (e.g., one or more thermometer that detect ambient temperature), humidity sensor components, pressure sensor components (e.g., barometer), acoustic sensor components (e.g., one or more microphones that detect background noise), proximity sensor components (e.g., infrared sensors that detect nearby objects), gas sensors (e.g., gas detection sensors to detect concentrations of hazardous gases for safety or to measure pollutants in the atmosphere), or other components that may provide indications, measurements, or signals corresponding to a surrounding physical environment. The position components **938** may include location sensor components (e.g., a GPS receiver component), altitude sensor components (e.g., altimeters or barometers that detect air pressure from which altitude may be derived), orientation sensor components (e.g., magnetometers), and the like.

(111) Communication may be implemented using a wide variety of technologies. The I/O components **918** may include communication components **940** operable to couple the machine **900** to a network **932** or devices **920** via coupling **924** and coupling **922**, respectively. For example, the communication components **940** may include a network interface component or other suitable device to interface with the network **932**. In further examples, communication components **940** may include wired communication components, wireless communication components, cellular communication components, near field communication (NFC) components, Bluetooth® components (e.g., Bluetooth® Low Energy), Wi-Fi® components, and other communication components to provide communication via other modalities. The devices **920** may be another machine or any of a wide variety of peripheral devices (e.g., a peripheral device coupled via a USB).

(112) Moreover, the communication components **940** may detect identifiers or include components operable to detect identifiers. For example, the communication components **940** may include radio frequency identification (RFID) tag reader components, NFC smart tag detection components, optical reader components (e.g., an optical sensor to detect one-dimensional bar codes such as Universal Product Code (UPC) bar code, multi-dimensional bar codes such as Quick Response (QR) code, Aztec code, Data Matrix, Dataglyph, MaxiCode, PDF417, Ultra Code, UCC RSS-2D bar code, and other optical codes), or acoustic detection components (e.g., microphones to identify tagged audio signals). In addition, a variety of information may be derived via the communication components **940** such as location via Internet Protocol (IP) geo-location, location via Wi-Fi® signal triangulation, location via detecting a NFC beacon signal that may indicate a particular location, and so forth.

Glossary

(113) “CARRIER SIGNAL” in this context refers to any intangible medium that is capable of storing, encoding, or carrying instructions **910** for execution by the machine **900**, and includes digital or analog communications signals or other intangible medium to facilitate communication of such instructions **910**. Instructions **910** may be transmitted or received over the network **932** using a transmission medium via a network interface device and using any one of a number of

well-known transfer protocols.

(114) “CLIENT DEVICE” in this context refers to any machine **900** that interfaces to a communications network **932** to obtain resources from one or more server systems or other client devices. A client device may be, but is not limited to, mobile phones, desktop computers, laptops, PDAs, smart phones, tablets, ultra books, netbooks, laptops, multi-processor systems, microprocessor-based or programmable consumer electronics, game consoles, STBs, or any other communication device that a user may use to access a network **932**.

(115) “COMMUNICATIONS NETWORK” in this context refers to one or more portions of a network **932** that may be an ad hoc network, an intranet, an extranet, a virtual private network (VPN), a LAN, a wireless LAN (WLAN), a WAN, a wireless WAN (WWAN), a metropolitan area network (MAN), the Internet, a portion of the Internet, a portion of the Public Switched Telephone Network (PSTN), a plain old telephone service (POTS) network, a cellular telephone network, a wireless network, a Wi-Fi® network, another type of network, or a combination of two or more such networks. For example, a network **932** or a portion of a network **932** may include a wireless or cellular network and the coupling may be a Code Division Multiple Access (CDMA) connection, a Global System for Mobile communications (GSM) connection, or other type of cellular or wireless coupling. In this example, the coupling may implement any of a variety of types of data transfer technology, such as Single Carrier Radio Transmission Technology (1×RTT), Evolution-Data Optimized (EVDO) technology, General Packet Radio Service (GPRS) technology, Enhanced Data rates for GSM Evolution (EDGE) technology, third Generation Partnership Project (3GPP) including 3G, fourth generation wireless (4G) networks, Universal Mobile Telecommunications System (UMTS), High Speed Packet Access (HSPA), Worldwide Interoperability for Microwave Access (WiMAX), Long Term Evolution (LTE) standard, others defined by various standard setting organizations, other long range protocols, or other data transfer technology.

(116) “MACHINE-READABLE MEDIUM” in this context refers to a component, device or other tangible media able to store instructions **910** and data temporarily or permanently and may include, but is not be limited to, random-access memory (RAM), read-only memory (ROM), buffer memory, flash memory, optical media, magnetic media, cache memory, other types of storage (e.g., erasable programmable read-only memory (EEPROM)), and/or any suitable combination thereof. The term “machine-readable medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database, or associated caches and servers) able to store instructions **910**. The term “machine-readable medium” shall also be taken to include any medium, or combination of multiple media, that is capable of storing instructions **910** (e.g., code) for execution by a machine **900**, such that the instructions **910**, when executed by one or more processors **904** of the machine **900**, cause the machine **900** to perform any one or more of the methodologies described herein. Accordingly, a “machine-readable medium” refers to a single storage apparatus or device, as well as “cloud-based” storage systems or storage networks that include multiple storage apparatus or devices. The term “machine-readable medium” excludes signals per se.

(117) “COMPONENT” in this context refers to a device, physical entity, or logic having boundaries defined by function or subroutine calls, branch points, APIs, or other technologies that provide for the partitioning or modularization of particular processing or control functions. Components may be combined via their interfaces with other components to carry out a machine process. A component may be a packaged functional hardware unit designed for use with other components and a part of a program that usually performs a particular function of related functions. Components may constitute either software components (e.g., code embodied on a machine-readable medium) or hardware components. A “hardware component” is a tangible unit capable of performing certain operations and may be configured or arranged in a certain physical manner. In various example embodiments, one or more computer systems (e.g., a standalone computer system,

a client computer system, or a server computer system) or one or more hardware components of a computer system (e.g., a processor or a group of processors **904**) may be configured by software (e.g., an application **816** or application portion) as a hardware component that operates to perform certain operations as described herein. A hardware component may also be implemented mechanically, electronically, or any suitable combination thereof. For example, a hardware component may include dedicated circuitry or logic that is permanently configured to perform certain operations. A hardware component may be a special-purpose processor, such as a field-programmable gate array (FPGA) or an application specific integrated circuit (ASIC). A hardware component may also include programmable logic or circuitry that is temporarily configured by software to perform certain operations. For example, a hardware component may include software executed by a general-purpose processor **904** or other programmable processor **904**. Once configured by such software, hardware components become specific machines **900** (or specific components of a machine **900**) uniquely tailored to perform the configured functions and are no longer general-purpose processors **904**. It will be appreciated that the decision to implement a hardware component mechanically, in dedicated and permanently configured circuitry, or in temporarily configured circuitry (e.g., configured by software), may be driven by cost and time considerations. Accordingly, the phrase “hardware component” (or “hardware-implemented component”) should be understood to encompass a tangible entity, be that an entity that is physically constructed, permanently configured (e.g., hardwired), or temporarily configured (e.g., programmed) to operate in a certain manner or to perform certain operations described herein. Considering embodiments in which hardware components are temporarily configured (e.g., programmed), each of the hardware components need not be configured or instantiated at any one instance in time. For example, where a hardware component comprises a general-purpose processor **904** configured by software to become a special-purpose processor, the general-purpose processor **904** may be configured as respectively different special-purpose processors (e.g., comprising different hardware components) at different times. Software accordingly configures a particular processor or processors **904**, for example, to constitute a particular hardware component at one instance of time and to constitute a different hardware component at a different instance of time. Hardware components can provide information to, and receive information from, other hardware components. Accordingly, the described hardware components may be regarded as being communicatively coupled. Where multiple hardware components exist contemporaneously, communications may be achieved through signal transmission (e.g., over appropriate circuits and buses **902**) between or among two or more of the hardware components. In embodiments in which multiple hardware components are configured or instantiated at different times, communications between such hardware components may be achieved, for example, through the storage and retrieval of information in memory structures to which the multiple hardware components have access. For example, one hardware component may perform an operation and store the output of that operation in a memory device to which it is communicatively coupled. A further hardware component may then, at a later time, access the memory device to retrieve and process the stored output. Hardware components may also initiate communications with input or output devices, and can operate on a resource (e.g., a collection of information). The various operations of example methods described herein may be performed, at least partially, by one or more processors **904** that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Whether temporarily or permanently configured, such processors **904** may constitute processor-implemented components that operate to perform one or more operations or functions described herein. As used herein, “processor-implemented component” refers to a hardware component implemented using one or more processors **904**. Similarly, the methods described herein may be at least partially processor-implemented, with a particular processor or processors **904** being an example of hardware. For example, at least some of the operations of a method may be performed by one or more processors **904** or processor-implemented components. Moreover, the

one or more processors **904** may also operate to support performance of the relevant operations in a “cloud computing” environment or as a “software as a service” (SaaS). For example, at least some of the operations may be performed by a group of computers (as examples of machines **900** including processors **904**), with these operations being accessible via a network **932** (e.g., the Internet) and via one or more appropriate interfaces (e.g., an API). The performance of certain of the operations may be distributed among the processors **904**, not only residing within a single machine **900**, but deployed across a number of machines **900**. In some example embodiments, the processors **904** or processor-implemented components may be located in a single geographic location (e.g., within a home environment, an office environment, or a server farm). In other example embodiments, the processors **904** or processor-implemented components may be distributed across a number of geographic locations.

(118) “PROCESSOR” in this context refers to any circuit or virtual circuit (a physical circuit emulated by logic executing on an actual processor **904**) that manipulates data values according to control signals (e.g., “commands,” “op codes,” “machine code,” etc.) and which produces corresponding output signals that are applied to operate a machine **900**. A processor **904** may be, for example, a central processing unit (CPU), a reduced instruction set computing (RISC) processor, a complex instruction set computing (CISC) processor, a graphics processing unit (GPU), a digital signal processor (DSP), an ASIC, a radio-frequency integrated circuit (RFIC) or any combination thereof. A processor **904** may further be a multi-core processor having two or more independent processors **904** (sometimes referred to as “cores”) that may execute instructions **910** contemporaneously.

Claims

1. A method comprising: receiving, by an input port of a packet forwarding system, a first data unit comprising a first header portion and a first payload portion; processing, by a parser of the packet forwarding system, the first data unit to generate a first parsing output that describes the first header portion; determining, by a header lookup functional component of the packet forwarding system, a forwarding decision for the first data unit based on the first parsing output, the determining of the forwarding decision comprising performing a three-way branch decision to either perform a header lookup operation based on the first parsing output, recirculate the first data unit back to the parser to generate another parsing output that describes another header portion, or provide an output of the lookup operation, the forwarding decision being determined based on the output of the header lookup operation; selecting, by the header lookup functional component of the packet forwarding system, an output queue for the first data unit based on the forwarding decision; modifying, by a header rewrite functional component of the packet forwarding system, a copy of the first header portion stored in an external buffer to generate a modified first data unit, the modifying of the copy of the first header portion comprising: removing, from the copy of the first header portion, an existing encapsulating header associated with a terminated tunnel; generating a new encapsulating header based on a template, the template providing one or more tunnel-specific values that are included in the new encapsulating header when the new encapsulating header is generated, the one or more tunnel-specific values comprising at least one of a source address or a destination address, and the one or more tunnel-specific values comprising a time-to-live value and a quality of service code; and adding, to the first header portion, the new encapsulating header; adding the modified first data unit to the output queue; and transmitting the modified first data unit via a transmit port of the packet forwarding system that is serviced by the output queue.
2. The method of claim 1, wherein the modifying of the copy of the first header portion comprises: updating a forwarding header in the copy of the first header portion based on at least one of a forwarding method being used to transmit the first data unit or a network path being used to transmit the first data unit.

3. The method of claim 1, wherein the output of the header lookup operation comprising either a metadata item, the metadata item and a header byte string, the metadata item and the first data unit in its entirety, or the metadata item and a portion of the first data unit that is greater than the header byte string.
4. A system comprising: a port manager functional component comprising an input port and a transmit port, wherein the port manager functional component is configured to perform operations comprising: receiving a first data unit via the input port, the first data unit comprising a first header portion and a first payload portion; and transmitting a modified first data unit via the transmit port that is serviced by an output queue; a parser configured to process the first data unit to generate a first parsing output that describes the first header portion; a header lookup functional component configured to perform operations comprising: determining a forwarding decision for the first data unit based on the first parsing output, the determining of the forwarding decision comprising performing a three-way branch decision to either perform a header lookup operation based on the first parsing output, recirculate the first data unit back to the parser to generate another parsing output that describes another header portion, or provide an output of the lookup operation, the forwarding decision being determined based on the output of the header lookup operation; and selecting an output queue for the first data unit based on the forwarding decision; and a header rewrite functional component configured to generate the modified first data unit and to add the modified first data unit to the output queue, the header rewrite functional component generating the modified first data unit by modifying a copy of the first header portion stored in an external buffer, the modifying of the copy of the first header portion comprising: removing, from the copy of the first header portion, an existing encapsulating header associated with a terminated tunnel; generating a new encapsulating header based on a template, the template providing one or more tunnel-specific values that are included in the new encapsulating header when the new encapsulating header is generated, the one or more tunnel-specific values comprising at least one of a source address or a destination address, and the one or more tunnel-specific values comprising at least one of a time-to-live value or a quality of service code; and adding, to the first header portion, the new encapsulating header.
5. The system of claim 4, wherein the modifying of the copy of the first header portion comprises: updating a forwarding header in the copy of the first header portion based on at least one of a forwarding method being used to transmit the first data unit or a network path being used to transmit the first data unit.
6. The system of claim 4, wherein the output of the header lookup operation comprising either a metadata item, the metadata item and a header byte string, the metadata item and the first data unit in its entirety, or the metadata item and a portion of the first data unit that is greater than the header byte string.
7. A physical device configured to perform operations comprising: receiving, by an input port of the physical device, a first data unit comprising a first header portion and a first payload portion; processing, by a parser of the physical device, the first data unit to generate a first parsing output that describes the first header portion; determining, by a header lookup functional component of the physical device, a forwarding decision for the first data unit based on the first parsing output, the determining of the forwarding decision comprising performing a three-way branch decision to either perform a header lookup operation based on the first parsing output, recirculate the first data unit back to the parser to generate another parsing output that describes another header portion, or provide an output of the lookup operation, the forwarding decision being determined based on the output of the header lookup operation; selecting, by the header lookup functional component of the physical device, an output queue for the first data unit based on the forwarding decision; modifying, by a header rewrite functional component of the physical device, a copy of the first header portion stored in an external buffer to generate a modified first data unit, the modifying of the copy of the first header portion comprising: removing, from the copy of the first header portion,

an existing encapsulating header associated with a terminated tunnel; generating a new encapsulating header based on a template, the template providing one or more tunnel-specific values that are included in the new encapsulating header when the new encapsulating header is generated, the one or more tunnel-specific values comprising at least one of a source address or a destination address, and the one or more tunnel-specific values comprising at least one of a time-to-live value or a quality of service code; and adding, to the first header portion, the new encapsulating header; adding the modified first data unit to the output queue; and transmitting the modified first data unit via a transmit port of the physical device that is serviced by the output queue.

8. The physical device of claim 7, wherein the output of the header lookup operation comprising either a metadata item, the metadata item and a header byte string, the metadata item and the first data unit in its entirety, or the metadata item and a portion of the first data unit that is greater than the header byte string.
