

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12393363
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Muchherla; Kishore Kumar et al.

---

### Voltage bin calibration based on a voltage distribution reference voltage

---

#### Abstract

An amount of voltage shift is determined for one or more memory cells of a block family based on an initial reference value pertaining to the one or more memory cells and a subsequent reference value pertaining to the one or more memory cells. The block family is associated with a first voltage bin or a second voltage bin based on the determined amount of voltage shift. The first voltage bin is associated with a first voltage offset and the second voltage bin is associated with a second voltage offset.

---

**Inventors:** Muchherla; Kishore Kumar (San Jose, CA), Batutis; Devin M. (San Jose, CA), Luo; Xiangang (Fremont, CA), Kaynak; Mustafa N. (San Diego, CA), Feeley; Peter (Boise, ID), Parthasarathy; Sivagnanam (Carlsbad, CA), Ratnam; Sampath (Boise, ID), Nowell; Shane (Boise, ID)

**Applicant:** MICRON TECHNOLOGY, INC. (Boise, ID)

**Family ID:** 1000008762855

**Assignee:** Micron Technology, Inc. (Boise, ID)

**Appl. No.:** 18/616006

**Filed:** March 25, 2024

#### Prior Publication Data

Document Identifier	Publication Date
US 20240231676 A1	Jul. 11, 2024

#### Related U.S. Application Data

continuation parent-doc US 18175439 20230227 US 11966616 child-doc US 18616006  
continuation parent-doc US 17203474 20210316 US 11620074 20230404 child-doc US 18175439

---

## Publication Classification

**Int. Cl.:** G06F3/00 (20060101); G06F3/06 (20060101)

**U.S. Cl.:**

**CPC** G06F3/0653 (20130101); G06F3/0604 (20130101); G06F3/0655 (20130101);  
G06F3/0679 (20130101);

## Field of Classification Search

**USPC:** None

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5610715	12/1996	Yoshii et al.	N/A	N/A
6377060	12/2001	Burkhart et al.	N/A	N/A
2017/0117053	12/2016	Sharon	N/A	G06F 11/1068
2019/0107961	12/2018	Lee	N/A	N/A
2020/0250030	12/2019	Sandell et al.	N/A	N/A
2022/0050777	12/2021	Sheperek et al.	N/A	N/A

---

*Primary Examiner:* Rojas; Midys

*Attorney, Agent or Firm:* Lowenstein Sandler LLP

---

## Background/Summary

RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 18/175,439, filed Feb. 27, 2023, which is a continuation of U.S. patent application Ser. No. 17/203,474, filed Mar. 16, 2021, now U.S. Pat. No. 11,620,074, the entire contents of which is hereby incorporated by reference herein.

### TECHNICAL FIELD

(1) Embodiments of the disclosure relate generally to memory sub-systems, and more specifically, relate to voltage bin calibration based on a voltage distribution reference voltage.

### BACKGROUND

(2) A memory sub-system can include one or more memory devices that store data. The memory devices can be, for example, non-volatile memory devices and volatile memory devices. In general, a host system can utilize a memory sub-system to store data at the memory devices and to retrieve data from the memory devices.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) The disclosure will be understood more fully from the detailed description given below and

from the accompanying drawings of various embodiments of the disclosure. The drawings, however, should not be taken to limit the disclosure to the specific embodiments, but are for explanation and understanding only.

(2) FIG. 1 illustrates an example computing system that includes a memory sub-system, in accordance with some embodiments of the present disclosure.

(3) FIG. 2 schematically illustrates the temporal voltage shift caused by the slow charge loss exhibited by triple-level memory cells, in accordance with some embodiments of the present disclosure.

(4) FIG. 3 depicts an example graph illustrating the dependency of the threshold voltage offset on the time after program (i.e., the period of time elapsed since the block had been programmed, in accordance with some embodiments of the present disclosure.

(5) FIG. 4 schematically illustrates a set of predefined threshold voltage offset bins, in accordance with embodiments of the present disclosure.

(6) FIG. 5 illustrates block family management operations implemented by a block family manager component, in accordance with embodiments of the present disclosure.

(7) FIG. 6 schematically illustrates selecting block families for calibration, in accordance with embodiments of the present disclosure.

(8) FIG. 7 is a flow diagram of an example method for voltage bin calibration based on a voltage distribution reference voltage, in accordance with some embodiments of the present disclosure.

(9) FIGS. 8A-8B illustrates an example of detecting voltage shift using a memory access voltage distribution reference voltage, in accordance with some embodiments of the present disclosure.

(10) FIGS. 9A-9B also illustrates another example of detecting voltage shift using a memory access voltage distribution reference voltage, in accordance with some embodiments of the present disclosure.

(11) FIG. 10 is a flow diagram of another example method for voltage bin calibration based on a voltage distribution reference voltage, in accordance with some embodiments of the present disclosure.

(12) FIG. 11 illustrates a change in the width of a memory access voltage distribution during operation of a memory device, in accordance with embodiments of the present disclosure.

(13) FIG. 12 is a block diagram of an example computer system in which embodiments of the present disclosure may operate.

#### DETAILED DESCRIPTION

(14) Aspects of the present disclosure are directed to voltage bin calibration based on a voltage distribution reference voltage. A memory sub-system can be a storage device, a memory module, or a combination of a storage device and memory module. Examples of storage devices and memory modules are described below in conjunction with FIG. 1. In general, a host system can utilize a memory sub-system that includes one or more memory components, such as memory devices that store data. The host system can provide data to be stored at the memory sub-system and can request data to be retrieved from the memory sub-system.

(15) A memory sub-system can utilize one or more memory devices, including any combination of the different types of non-volatile memory devices and/or volatile memory devices, to store the data provided by the host system. In some embodiments, non-volatile memory devices can be provided by negative-and (NAND) type flash memory devices. Other examples of non-volatile memory devices are described below in conjunction with FIG. 1. A non-volatile memory device is a package of one or more dice. Each die can consist of one or more planes. Planes can be groups into logic units (LUN). For some types of non-volatile memory devices (e.g., NAND devices), each plane consists of a set of physical blocks. Each block consists of a set of pages. “Block” herein shall refer to a set of contiguous or non-contiguous memory pages. An example of a “block” is an “erasable block,” which is the minimal erasable unit of memory, while “page” is a minimal writable unit of memory. Each page includes a set of memory cells. A memory cell is an electronic

circuit that stores information.

(16) Data operations can be performed by the memory sub-system. The data operations can be host-initiated operations. For example, the host system can initiate a data operation (e.g., write, read, erase, etc.) on a memory sub-system. The host system can send access requests (e.g., write command, read command) to the memory sub-system, such as to store data on a memory device at the memory sub-system and to read data from the memory device on the memory sub-system. The data to be read or written, as specified by a host request, is hereinafter referred to as “host data.” A host request can include logical address information (e.g., logical block address (LBA), namespace) for the host data, which is the location the host system associates with the host data. The logical address information (e.g., LBA, namespace) can be part of metadata for the host data. Metadata can also include error handling data (e.g., ECC codeword, parity code), data version (e.g. used to distinguish age of data written), valid bitmap (which LBAs or logical transfer units contain valid data), etc.

(17) A memory device includes multiple memory cells, each of which can store, depending on the memory cell type, one or more bits of information. A memory cell can be programmed (written to) by applying a certain voltage to the memory cell, which results in an electric charge being held by the memory cell, thus allowing modulation of the voltage distributions produced by the memory cell. Moreover, precisely controlling the amount of the electric charge stored by the memory cell allows to establish multiple threshold voltage levels corresponding to different logical levels, thus effectively allowing a single memory cell to store multiple bits of information: a memory cell operated with  $2^{\text{sup.}n}$  different threshold voltage levels is capable of storing  $n$  bits of information. “Threshold voltage” herein shall refer to the voltage level that defines a boundary between two neighboring voltage distributions corresponding to two logical levels. Thus, a read operation can be performed by comparing the measured voltage exhibited by the memory cell to one or more reference voltage levels in order to distinguish between two logical levels for single-level cells and between multiple logical levels for multi-level cells.

(18) Due to the phenomenon known as slow charge loss, the threshold voltage of a memory cell changes in time as the electric charge of the cell is degrading, which is referred to as “temporal voltage shift” (since the degrading electric charge causes the voltage distributions to shift along the voltage axis towards lower voltage levels). The threshold voltage is changing rapidly at first (immediately after the memory cell was programmed), and then slows down in an approximately logarithmic linear fashion with respect to the time elapsed since the cell programming event. Accordingly, failure to mitigate the temporal voltage shift caused by the slow charge loss can result in the increased bit error rate in read operations.

(19) However, various common implementations either fail to adequately address the temporal voltage shift or employ inefficient strategies resulting in high bit error rates and/or exhibiting other shortcomings. Embodiments of the present disclosure address the above-noted and other deficiencies by implementing a memory sub-system that employs block family based error avoidance strategies, thus significantly improving the bit error rate exhibited by the memory sub-system.

(20) In accordance with embodiments of the present disclosure, the temporal voltage shift is selectively tracked for a programmed set of memory cells grouped by block families, and appropriate voltage offsets, which are based on block affiliation with a certain block family, are applied to the base read levels in order to perform read operations. “Block family” herein shall refer to a possibly noncontiguous set of memory cells (which can reside in one or more full and/or partial blocks, the latter referred to as “partitions” herein) that have been programmed within a specified time window and a specified temperature window, and thus are expected to exhibit similar or correlated changes in their respective data state metrics for slow charge loss. A block family can be made with any granularity containing only whole codewords, whole pages, whole super pages, or whole superblocs, or any combination of these. “Data state metric” herein shall

refer to a quantity that is measured or inferred from the state of data stored on a memory device. Specifically, the data state metrics can reflect the state of the temporal voltage shift, the degree of read disturb, and/or other measurable functions of the data state. A composite data state metric is a function (e.g., a weighted sum) of a set of component state metrics.

(21) In these embodiments, a block family experiences a specified temperature range using an aggregate temperature, which can change or be defined according to a temperature metric in different embodiments. For example, in response to a timer reaching a soft closure value, a controller of the memory sub-system can perform a soft closure of the block family (i.e., signal to a block family manager that the open time for the block is coming to an end and prepare for a hard closure of the block family). Or, alternatively, in response to an aggregate temperature (based on temperature values received from temperature sensor(s) of the memory device) being greater than or equal to a specified threshold temperature, the controller can likewise perform a soft closure of the block family.

(22) After performing the soft closure, the memory sub-system can continue to program to a partially written block until the timer reaches a hard closure value or the block family satisfies a hard closure criteria. In this way, the memory sub-system has an extension time before the hard closure during which to complete writing to the partially written block, decide whether to partition the block, or to write dummy data to the unwritten portion of the block after hard closure. After hard closure of one block family (i.e., the memory sub-system controller no longer writes data to the block family), further programming to dice of the memory device is to a newly opened block family. Given that wear-leveling keeps block to similar program erase counts, the time elapsed after programming and temperature are the main factors affecting the temporal voltage shift, all blocks and/or partitions within a single block family are presumed to exhibit similar distributions of threshold voltages in memory cells, and thus would require the same voltage offsets to be applied to the base read levels for read operations. “Base read level” herein shall refer to the initial threshold voltage level exhibited by the memory cell immediately after programming. In some implementations, base read levels can be stored in the metadata of the memory device.

(23) Block families can be created asynchronously with respect to block programming events. In an illustrative example, a new block family can be created whenever a specified period of time (e.g., a predetermined number of minutes) has elapsed since creation of the last block family or the reference temperature of memory cells has changed by more than a specified threshold value. The memory sub-system controller can maintain an identifier of the active block family, which is associated with one or more blocks as they are being programmed.

(24) The memory sub-system controller can periodically perform a calibration process (e.g., a calibration scan) in order to associate each die of every block family with one of the predefined threshold voltage offset bins (referred to as voltage bins), which is in turn associated with the voltage offsets to be applied for read operations. In one embodiment, the associations of pages or blocks with block families, and block families and dies with voltage bins, can be represented by metadata stored in respective metadata tables maintained by the memory sub-system controller.

(25) The threshold voltage for each memory cell can continue to shift even when a memory device is powered off. Depending on the amount of time the memory device is powered off, the temporal voltage shift of each block family can be significant. As a result, the bin allocation to block families stored in respective metadata tables prior to the memory device being powered off can be inaccurate when the memory device is powered back on. In some instances, a memory sub-system does not have a timer or a clock that tracks the amount of time a memory device is powered off. As a result, when the memory device is powered back on, the memory sub-system controller is unaware of the amount of time the device was powered off and is therefore unaware of the temporal voltage shift that has occurred for each block family. Thus, the memory sub-system controller cannot adjust the voltage bin allocation to block families without calibration of at least some pages or blocks.

(26) The memory sub-system controller can periodically perform a scan of the memory device to calibrate the threshold voltage offsets to be applied to each block family (referred to as a calibration scan). In some instances, the memory sub-system controller can perform the calibration scan by performing a set of read operations performed for a block family using voltage offsets associated with each voltage bin of the memory device. The memory sub-system controller can identify a previously performed read operation that resulted in an optimal data state metric (e.g., the lowest bit error rate, which is the number of received bits that have been altered due to an error at the memory sub-system) and determine the voltage offset that has been applied for the identified read operation. In one example, the memory sub-system controller can perform the calibration scan for each block of each block family on each die of the memory device. As a memory device can include hundreds or sometimes thousands of pages or blocks, periodically performing a full scan of the memory device can take a significant amount of time and utilize a significant amount of system resources.

(27) Aspects of the present disclosure address the above and other deficiencies by calibrating voltage bins at a memory device based on a voltage shift detected using a reference voltage for a memory access voltage distribution (i.e., a threshold voltage distribution). The reference voltage can correspond to a particular point (e.g., a tail end) or a particular quantile (e.g., a 50% quantile) of a voltage distribution for a respective logical level of memory cells at a memory device. At a time period immediately or soon after programming a block family, the memory pages of the block family can be associated with a particular memory access voltage distribution (referred to as a prior memory access voltage distribution). In one example, the reference voltage can correspond to a voltage for memory cells at a tail end of the voltage distribution. A memory access voltage for memory cells associated with the tail end of the prior memory access voltage distribution can approximately correspond to the threshold voltage used by the memory sub-system controller to program data to the memory cells (referred to as the programming voltage). As such, a value of the reference voltage that is associated with the prior memory access voltage distribution (referred to as the prior value of the reference voltage) corresponds to the programming voltage for the memory cells. In another example, the reference voltage can correspond to a voltage for memory cells at a particular quantile (e.g., a 50% quantile) of the voltage distribution.

(28) During operation of a memory device, a memory sub-system controller can measure a memory access voltage distribution for memory cells of a memory page of a block family (referred to as a current memory access voltage distribution). The memory sub-system controller can determine a current value for the reference voltage based on the current memory access voltage distribution and compare the current value to the prior value for the reference voltage. For example, the memory sub-system controller can determine a difference between the current voltage for memory cells associated with the particular point (e.g., the tail end) or the particular quantile (e.g., the 50% quantile) of the current memory access voltage distribution and the prior value for the memory cells associated with the corresponding point or the corresponding quantile of the prior memory access voltage distribution. The difference between the current value and the prior value for the reference voltage correspond to an amount of voltage shift of the block family since the time period immediately or soon after the memory sub-system controller programmed memory pages of the block family. The voltage shift can be dependent on an amount of time that has passed since the memory sub-system controller programmed the memory pages, a change in temperature from the time of programming and/or a number of program/erase (PE) cycles for the memory. The memory sub-system controlled can associate the block family with an appropriate voltage bin that corresponds with the determined amount of voltage shift.

(29) Advantages of the present disclosure include, but are not limited to, decreasing an amount of memory sub-system resources consumed during a calibration scan for the memory device. Instead of performing read operations at hundreds or thousands of pages or blocks to determine an appropriate voltage offset for each block family, the memory sub-system controller can detect the

amount of voltage shift for memory pages of a block family based on a difference between a current value and a prior value for a reference voltage associated with a voltage distribution measured for a memory page. By performing a single measurement operation for a memory page instead of multiple read operations for hundreds or thousands of pages or blocks, a significant amount of memory sub-system resources are available for other processes at the memory sub-system. As a result, an overall efficiency of the memory sub-system increases and an overall latency of the memory sub-system decreases. Additionally, by using the programming voltage for memory cells to determine the prior value of a reference voltage, the memory sub-system controller does not measure the voltage distribution of the memory page immediately or soon after programming. Memory sub-system resources can therefore be allocated to performing memory access operations with respect to host data, which can increase overall performance (e.g., efficiency, latency, throughput, etc.) of the memory sub-system.

(30) FIG. 1 illustrates an example computing system **100** that includes a memory sub-system **110** in accordance with some embodiments of the present disclosure. The memory sub-system **110** can include media, such as one or more volatile memory devices (e.g., memory device **140**), one or more non-volatile memory devices (e.g., memory device **130**), or a combination of such.

(31) A memory sub-system **110** can be a storage device, a memory module, or a combination of a storage device and memory module. Examples of a storage device include a solid-state drive (SSD), a flash drive, a universal serial bus (USB) flash drive, an embedded Multi-Media Controller (eMMC) drive, a Universal Flash Storage (UFS) drive, a secure digital (SD) card, and a hard disk drive (HDD). Examples of memory modules include a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), and various types of non-volatile dual in-line memory modules (NVDIMMs).

(32) The computing system **100** can be a computing device such as a desktop computer, laptop computer, network server, mobile device, a vehicle (e.g., airplane, drone, train, automobile, or other conveyance), Internet of Things (IoT) enabled device, embedded computer (e.g., one included in a vehicle, industrial equipment, or a networked commercial device), or such computing device that includes memory and a processing device.

(33) The computing system **100** can include a host system **120** that is coupled to one or more memory sub-systems **110**. In some embodiments, the host system **120** is coupled to multiple memory sub-systems **110** of different types. FIG. 1 illustrates one example of a host system **120** coupled to one memory sub-system **110**. As used herein, “coupled to” or “coupled with” generally refers to a connection between components, which can be an indirect communicative connection or direct communicative connection (e.g., without intervening components), whether wired or wireless, including connections such as electrical, optical, magnetic, etc.

(34) The host system **120** can include a processor chipset and a software stack executed by the processor chipset. The processor chipset can include one or more cores, one or more caches, a memory controller (e.g., NVDIMM controller), and a storage protocol controller (e.g., PCIe controller, SATA controller). The host system **120** uses the memory sub-system **110**, for example, to write data to the memory sub-system **110** and read data from the memory sub-system **110**.

(35) The host system **120** can be coupled to the memory sub-system **110** via a physical host interface. Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, universal serial bus (USB) interface, Fibre Channel, Serial Attached SCSI (SAS), a double data rate (DDR) memory bus, Small Computer System Interface (SCSI), a dual in-line memory module (DIMM) interface (e.g., DIMM socket interface that supports Double Data Rate (DDR)), etc. The physical host interface can be used to transmit data between the host system **120** and the memory sub-system **110**. The host system **120** can further utilize an NVM Express (NVMe) interface to access components (e.g., memory devices **130**) when the memory sub-system **110** is coupled with the host system **120** by the physical host interface (e.g., PCIe bus). The physical host

interface can provide an interface for passing control, address, data, and other signals between the memory sub-system **110** and the host system **120**. FIG. **1** illustrates a memory sub-system **110** as an example. In general, the host system **120** can access multiple memory sub-systems via a same communication connection, multiple separate communication connections, and/or a combination of communication connections.

(36) The memory devices **130**, **140** can include any combination of the different types of non-volatile memory devices and/or volatile memory devices. The volatile memory devices (e.g., memory device **140**) can be, but are not limited to, random access memory (RAM), such as dynamic random access memory (DRAM) and synchronous dynamic random access memory (SDRAM).

(37) Some examples of non-volatile memory devices (e.g., memory device **130**) include a negative-and (NAND) type flash memory and write-in-place memory, such as a three-dimensional cross-point (“3D cross-point”) memory device, which is a cross-point array of non-volatile memory cells. A cross-point array of non-volatile memory cells can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many flash-based memories, cross-point non-volatile memory can perform a write in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased. NAND type flash memory includes, for example, two-dimensional NAND (2D NAND) and three-dimensional NAND (3D NAND).

(38) Each of the memory devices **130** can include one or more arrays of memory cells. One type of memory cell, for example, single level cells (SLC) can store one bit per cell. Other types of memory cells, such as multi-level cells (MLCs), triple level cells (TLCs), quad-level cells (QLCs), and penta-level cells (PLCs) can store multiple bits per cell. In some embodiments, each of the memory devices **130** can include one or more arrays of memory cells such as SLCs, MLCs, TLCs, QLCs, PLCs or any combination of such. In some embodiments, a particular memory device can include an SLC portion, and an MLC portion, a TLC portion, a QLC portion, or a PLC portion of memory cells. The memory cells of the memory devices **130** can be grouped as pages that can refer to a logical unit of the memory device used to store data. With some types of memory (e.g., NAND), pages can be grouped to form blocks.

(39) Although non-volatile memory components such as a 3D cross-point array of non-volatile memory cells and NAND type flash memory (e.g., 2D NAND, 3D NAND) are described, the memory device **130** can be based on any other type of non-volatile memory, such as read-only memory (ROM), phase change memory (PCM), self-selecting memory, other chalcogenide based memories, ferroelectric transistor random-access memory (FeTRAM), ferroelectric random access memory (FeRAM), magneto random access memory (MRAM), Spin Transfer Torque (STT)-MRAM, conductive bridging RAM (CBRAM), resistive random access memory (RRAM), oxide based RRAM (OxRAM), negative-or (NOR) flash memory, or electrically erasable programmable read-only memory (EEPROM).

(40) A memory sub-system controller **115** (or controller **115** for simplicity) can communicate with the memory devices **130** to perform operations such as reading data, writing data, or erasing data at the memory devices **130** and other such operations. The memory sub-system controller **115** can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, or a combination thereof. The hardware can include a digital circuitry with dedicated (i.e., hard-coded) logic to perform the operations described herein. The memory sub-system controller **115** can be a microcontroller, special purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), etc.), or other suitable processor.

(41) The memory sub-system controller **115** can include a processing device, which includes one or more processors (e.g., processor **117**), configured to execute instructions stored in a local memory **119**. In the illustrated example, the local memory **119** of the memory sub-system controller **115** includes an embedded memory configured to store instructions for performing various processes,



operations, logic flows, and routines that control operation of the memory sub-system **110**, including handling communications between the memory sub-system **110** and the host system **120**. (42) In some embodiments, the local memory **119** can include memory registers storing memory pointers, fetched data, etc. The local memory **119** can also include read-only memory (ROM) for storing micro-code. While the example memory sub-system **110** in FIG. **1** has been illustrated as including the memory sub-system controller **115**, in another embodiment of the present disclosure, a memory sub-system **110** does not include a memory sub-system controller **115**, and can instead rely upon external control (e.g., provided by an external host, or by a processor or controller separate from the memory sub-system).

(43) In general, the memory sub-system controller **115** can receive commands or operations from the host system **120** and can convert the commands or operations into instructions or appropriate commands to achieve the desired access to the memory devices **130**. The memory sub-system controller **115** can be responsible for other operations such as wear leveling operations, garbage collection operations, error detection and error-correcting code (ECC) operations, encryption operations, caching operations, and address translations between a logical address (e.g., a logical block address (LBA), namespace) and a physical address (e.g., physical block address) that are associated with the memory devices **130**. The memory sub-system controller **115** can further include host interface circuitry to communicate with the host system **120** via the physical host interface. The host interface circuitry can convert the commands received from the host system into command instructions to access the memory devices **130** as well as convert responses associated with the memory devices **130** into information for the host system **120**.

(44) The memory sub-system **110** can also include additional circuitry or components that are not illustrated. In some embodiments, the memory sub-system **110** can include a cache or buffer (e.g., DRAM) and address circuitry (e.g., a row decoder and a column decoder) that can receive an address from the memory sub-system controller **115** and decode the address to access the memory devices **130**.

(45) In some embodiments, the memory devices **130** include local media controllers **135** that operate in conjunction with memory sub-system controller **115** to execute operations on one or more memory cells of the memory devices **130**. An external controller (e.g., memory sub-system controller **115**) can externally manage the memory device **130** (e.g., perform media management operations on the memory device **130**). In some embodiments, memory sub-system **110** is a managed memory device, which is a raw memory device **130** having control logic (e.g., local controller **132**) on the die and a controller (e.g., memory sub-system controller **115**) for media management within the same memory device package. An example of a managed memory device is a managed NAND (MNAND) device.

(46) The memory sub-system **110** includes a block family manager component **113** that can select threshold voltage offset bins (referred to as voltage bins or bins) to be associated with block families at a memory device. In some embodiments, the memory sub-system controller **115** includes at least a portion of the block family manager component **113**. For example, the memory sub-system controller **115** can include a processor **117** (processing device) configured to execute instructions stored in local memory **119** for performing the operations described herein. In some embodiments, the block family manager component **113** is part of the host system **110**, an application, or an operating system. Further details regarding block families and block family manager component **113** are described below.

(47) FIG. **2** schematically illustrates the temporal voltage shift caused by the slow charge loss exhibited by triple-level memory cells, in accordance with some embodiments of the present disclosure. While the illustrative example of FIG. **2** utilizes triple-level cells, the same observations can be made and, accordingly, the same remedial measures are applicable to single level cells and multi-level cells, as well as any other fractional or whole number of bits per cell (e.g., 3.5 bits per cell, etc.), in order to compensate for the slow charge loss.

(48) As noted above, a memory cell can be programmed (written to) by applying a certain voltage to the memory cell, which results in an electric charge being held by the memory cell, thus allowing modulation of the voltage distributions produced by the memory cell. Precisely controlling the amount of the electric charge stored by the memory cell allows to establish multiple threshold voltage levels corresponding to different logical levels, thus effectively allowing a single memory cell to store multiple bits of information: a memory cell operated with  $2^{\text{sup.}n}$  different threshold voltage levels is capable of storing  $n$  bits of information.

(49) In FIG. 2, each graph **220A-220N** shows a voltage distribution produced by memory cells programmed by a respective write level (which can be assumed to be at the midpoint of the distribution) to encode a corresponding logical level (e.g., “000” through “111” in case of a TLC). In order to distinguish between neighboring distributions (corresponding to two different logical levels), the read threshold voltage levels (shown by dashed vertical lines) are defined, such that any measured voltage that falls below a threshold level is associated with one distribution of the pair of neighboring distributions, while any measured voltage that is greater than or equal to the threshold level is associated with another distribution of the pair of neighboring distributions.

(50) As seen from comparing example charts **210** and **230**, which reflect the time periods immediately after programming and a respective number of hours after programming, respectively, the voltage distributions change in time due to the slow charge loss, which results in drifting values of the read threshold voltage levels, which are shown by dashed vertical lines. In various embodiments of the present disclosure, the temporal voltage shift is selectively tracked for programmed pages or blocks grouped by block families, and appropriate voltage offsets, which are based on page or block affiliation with a certain block family, are applied to the base read levels in order to perform read operations.

(51) FIG. 3 depicts an example graph illustrating the dependency of the threshold voltage offset on the time after program (i.e., the period of time elapsed since the page had been programmed, in accordance with some embodiments of the present disclosure. As schematically illustrated by FIG. 3, block families of the memory device are grouped into bins **330A-330N**, such that each block family includes one or more pages or blocks that have been programmed within a specified time window and a specified temperature window. As noted herein above, since the time elapsed after programming and temperature are the main factors affecting the temporal voltage shift, all pages, blocks, and/or partitions within a single block family **310** are presumed to exhibit similar distributions of threshold voltages in memory cells, and thus would require the same voltage changes for read operations.

(52) Block families can be created asynchronously with respect to page programming events. In an illustrative example, the memory sub-system controller **115** of FIG. 1 can create a new block family whenever a specified period of time (e.g., a predetermined number of minutes) has elapsed since creation of the last block family or whenever the reference temperature of memory cells, which is updated at specified time intervals, has changed by more than a specified threshold value since creation of the current block family.

(53) A newly created block family can be associated with a particular bin (e.g., bin 0, bin 1, etc.). Then, the memory sub-system controller can periodically perform a calibration process in order to associate each die of every block family with one of the predefined voltage bins (bins 0-7 in the illustrative example of FIG. 3), which is in turn associated with the voltage offset to be applied for read operations. The associations of blocks with block families and block families and dies with voltage bins can be stored in respective metadata tables maintained by the memory sub-system controller.

(54) FIG. 4 schematically illustrates a set of predefined threshold voltage offset bins (referred to as voltage bins or bins herein), in accordance with embodiments of the present disclosure. As schematically illustrated by FIG. 4, the threshold voltage offset graph can be subdivided into multiple voltage bins, such that each voltage bin corresponds to a predetermined range of threshold

voltage offsets. While the illustrative example of FIG. 4 defines ten voltage bins for a single valley, in other implementations, various other numbers of voltage bins can be employed.

(55) The memory sub-system controller can associate each die of every block family with a voltage bin, based on a periodically performed calibration process. The calibration process defines a set of threshold voltage offsets to be applied to the base voltage read level in order to perform read operations. In some embodiments, the calibration process involves performing, with respect to a specified number of randomly selected pages or blocks within the block family that is being calibrated, read operations utilizing different threshold voltage offsets, and choosing the threshold voltage offset that minimizes the error rate (e.g., a bit error rate) of the read operation. In other or similar embodiments, the calibration process involves sampling memory pages (i.e., measuring an access voltage distribution for memory cells of a memory page) of a block family and determining an amount voltage distribution shift from an access voltage distribution for the block family immediately or soon after programming. Further details regarding the calibration process are provided with respect to FIGS. 7 and 10.

(56) FIG. 5 schematically illustrates block family management operations implemented by the block family manager component 113 of the memory sub-system controller 115, in accordance with embodiments of the present disclosure. As schematically illustrated by FIG. 5, the block family manager component 113 can maintain, in a memory variable, an identifier 520 of the active block family, which is associated with one or more pages or blocks of cursors 530A-530K as they are being programmed. “Cursor” herein shall broadly refer to a location on the memory device to which the data is being written.

(57) The memory sub-system controller can utilize a power on minutes (POM) clock for tracking the creation times of block families. In some implementations, a less accurate clock, which continues running when the controller is in various low-power states, can be utilized in addition to the POM clock, such that the POM clock is updated based on the less accurate clock upon the controller wake-up from the low-power state.

(58) Thus, upon initialization of each block family, block family manager component 113 stores the current time 540 in a memory variable as the block family start time 550. As the pages or blocks are programmed, block family manager component 113 compares the current time 540 to the block family start time 550. Responsive to detecting that the difference of the current time 540 and the block family start time 550 is greater than or equal to the specified time period (e.g., a predetermined number of minutes), block family manager component 113 updates the memory variable storing the active block family identifier 520 to store the next block family number (e.g., the next sequential integer number), and the memory variable storing the block family start time 550 is updated to store the current time 540.

(59) The block family manager component 113 also maintains two memory variables for storing the high and low reference temperatures of a selected die of each memory device. Upon initialization of each block family, the high temperature 560 and the low temperature 570 variable store the value of the current temperature of the selected die of the memory device. In operation, while the active block family identifier 520 remains the same, temperature measurements are periodically obtained and compared with the stored high temperature 560 and the low temperature 570 values, which are updated accordingly: should the temperature measurement be found to be greater than or equal to the value stored by the high temperature variable 560, the latter is updated to store that temperature measurement; conversely, should the temperature measurement be found to fall below the value stored by the low temperature variable 570, the latter is updated to store that temperature measurement.

(60) The block family manager component 113 can further periodically compute the difference between the high temperature 560 and the low temperature 570. Responsive to determining that the difference between the high temperature 560 and the low temperature 570 is greater than or equal to a specified temperature threshold, the block family manager component 113 can close the

existing block family and to create a new active block family: the memory variable storing the active block family identifier **520** is updated to store the next block family number (e.g., the next sequential integer number), the memory variable storing the block family start time **550** is updated to store the current time **540**, and the high temperature **560** and the low temperature **570** variables are updated to store the value of the current temperature of the selected die of the memory device. At the time of programming a block, block family manager component **113** associates the block with the currently active block family. The association of each block with a corresponding block family is reflected by the block family metadata **580**.

(61) As described previously, based on a periodically performed calibration process (e.g., a calibration scan), the block family manager component **113** associates each die of every block family with a voltage bin, which defines a set of threshold voltage offsets to be applied to the base voltage read level in order to perform read operations. In some embodiments, the calibration process involves sampling memory cells for a memory page of a block family to measure a current memory access voltage distribution for the memory page (referred to as current voltage distribution herein). In some embodiments, the current voltage distribution can correspond to an average distribution for a set of sampled memory cells. In other or similar embodiments, the current voltage distribution can be a different aggregate measure. Block family manager component **113** can determine a current value for a reference voltage based on the current voltage distribution and compare the current value for the reference voltage to a prior value for the reference voltage. The prior value for the reference voltage is associated with a prior memory access voltage distribution (referred to as a prior voltage distribution) that associated with the block family at an instance in time that is prior to measuring the current memory access voltage distribution for the memory page (e.g., when data for the memory page is programmed at the memory device). Block family manager component **113** uses the current value and the prior value for the reference voltage to determine the amount of voltage shift for the block family since the instance in time associated with the prior reference value. The determined amount of voltage shift can correspond to a threshold voltage offset that is previously determined (e.g., based on experimental data, etc.) to minimize the error rate of operations such as read operations performed for the block family. Block family manager **113** can determine the particular voltage bin that corresponds to the threshold voltage offset and update metadata for the block family to correspond to the determined voltage bin.

(62) In some embodiments, the frequency at which the memory sub-system controller performs the calibration process for each voltage bin can be based on an age of the block families associated with the voltage bin. As described previously with respect to FIG. 3, newly created block families can be associated with voltage bin 0 and older block families on the memory device can be associated with subsequently numbered voltage bins. The temporal voltage shift for block families in a younger voltage bin is faster than the temporal voltage shift for block families associated with an older voltage bin. This is illustrated in FIG. 4, as the voltage offset for bin 0 shifts at quicker rate than the voltage offset for older voltage bins (e.g., voltage bins 9, 8, 7, etc.). Therefore, the memory sub-system controller can perform the calibration process for block families associated with voltage bin 0 at a higher frequency than for block families associated with voltage bin 9 to associate each block family with an appropriate voltage bin.

(63) FIG. 6 schematically illustrates selecting block families for calibration, in accordance with embodiments of the present disclosure. Due to slow charge loss, the oldest block families in a voltage bin will migrate to the next voltage bin before any other block families of the current bin. As such, the memory sub-system controller can limit calibration operations to the oldest block families in a bin (e.g., block family **610** in bin 0 and block family **620** in bin 1). In some embodiments, the memory sub-system controller can identify the oldest block family in a voltage bin based on a bin boundary for the bin. A bin boundary can represent a boundary between two adjacent block families that are each associated with a different bin. The memory sub-system controller can identify the bin boundary for a particular voltage bin using a block family metadata

table.

(64) Referring back to FIG. 1, block family manager component **113** can be configured to periodically calibrate a memory device **130, 140** to associate block families with an appropriate voltage bin in view of an amount of temporal voltage shift. The block family can be associated with a first voltage distribution at a first time period (e.g., immediately or soon after data is programmed or written to memory pages of the block family). In accordance with some embodiments of the present disclosure, the first voltage distribution is referred to as a prior memory access voltage distribution. In some embodiments, block family manager component **113** can obtain the first voltage distribution by measuring one or more memory pages of the block family during the first period. In other or similar embodiments, block family manager component **113** can determine the first voltage distribution based on experimental or manufacturing data that is obtained for the memory device (e.g., prior to or during installation and/or initialization of the memory device at the memory sub-system). A first value for a reference voltage (also referred to herein as a prior value) can be determined based on the first voltage distribution. A reference voltage refers to a voltage at a particular point (e.g., a tail end) or a particular quantile (e.g., 50% quantile, 25% quantile, etc.) of a voltage distribution. In some embodiments, the voltage at a quantile that is associated with a highest amount of slow charge loss can be selected as the reference voltage.

(65) At a second period of time subsequent to the first period of time, block family manager component **113** can perform a calibration scan for the block family by measuring one or more memory pages to obtain a second voltage distribution. In accordance with some embodiments of the present disclosure, the second voltage distribution is referred to as a current voltage distribution. Block family manager component **113** can compare a second value for the reference voltage (also referred to herein as a current value) to the first value for the reference voltage to determine the amount of temporal voltage shift for the block family since the first period of time. Block family manager component **113** can associate the block family with an appropriate voltage bin associated with a temporal voltage offset that corresponds to the determined amount of temporal voltage shift. Further details regarding the calibration process are provided with respect to FIGS. 7-11 herein.

(66) FIG. 7 is a flow diagram of an example method **700** for voltage bin calibration based on a voltage distribution reference voltage, in accordance with some embodiments of the present disclosure. The method **700** can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **700** is performed by the block family manager component **113** of FIG. 1. In other or similar embodiments, one or more operations of method **700** is performed by another component of the memory sub-system controller, such as a block family manager. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other process flows are possible.

(67) At block **710**, processing logic measures a current voltage distribution for a memory page of a block family. In some embodiments, the block family is associated with a first voltage bin associated with a first voltage offset at the memory device. The processing logic can measure the current voltage distribution for the memory page in response to determining that the block family is to be calibrated, in accordance with some embodiments. For example, the processing logic can determine to calibrate the block family in response to determining that a corresponding threshold criterion has been satisfied. In one embodiment, the processing logic determines that the threshold criterion is satisfied when a threshold number of program and erase operations have been

performed (i.e. a threshold number of program/erase (P/E) cycles have occurred) for the block family and/or the memory device. In some embodiments, the processing logic can select the memory page of the block family for measurement in accordance with a calibration policy associated with the memory device. In other or similar embodiments, the processing logic can select the memory page of the block family for measurement as a result of a random memory page selection operation performed during the calibration process.

(68) Processing logic can measure the current voltage distribution for the memory page by sampling the memory page to determine the memory access voltage associated with one or more cells of the memory page. In some embodiments, processing logic can measure the voltage distribution by sampling each memory cell of the memory page. In other or similar embodiments, processing logic can measure the voltage distribution by sampling a subset of memory cells of the memory page (i.e., fewer than all memory cells of the memory page). The current voltage distribution can represent a difference between the memory access voltage for the cells of the memory page at the time period that the measurement is obtained for the block family.

(69) FIGS. **8A** and **9A** illustrate current voltage distributions **810**, **820** measured for a memory page, in accordance with embodiments of the present disclosure. As described with respect to FIG. **2**, memory cells can be programmed by a respective write level to encode a corresponding logical level. In one example, voltage distributions **810** and **820** can be measured during a sampling operation performed for TLC memory cells of a memory page. In such example, voltage distribution **820** can correspond to a voltage distribution for a highest logical level associated with the TLC memory cell (e.g., “111”) and voltage distribution **810** can correspond to a voltage distribution for a second highest logical level associated with the TLC memory cell (e.g., “110”).

(70) Referring back to FIG. **7**, at block **720**, processing logic determines a current value for a reference voltage based on the current voltage distribution measured for the memory page. As described above, a reference voltage refers to a voltage at a particular point (e.g., a tail end of the distribution) or a particular quantile (e.g., 50% quantile, 25% quantile, etc.) of the voltage distribution. In some embodiments, the reference voltage can be defined for the memory device prior to or during initialization of the memory device (e.g., by an operator or programmer of the memory device). In other or similar embodiments, the reference voltage can be selected in view of a performance metric (e.g., bit error rate) associated with the block family and/or the memory device, as described in further details below.

(71) As described above, in some embodiments, the reference voltage can correspond to a voltage at particular point of the voltage distribution. For example, the reference voltage can correspond to a tail end of the voltage distribution. As illustrated in FIG. **8A**, voltage **812** can correspond to a voltage measured at a tail end of voltage distribution **810** and voltage **822** can correspond to a voltage measured at a tail end of voltage distribution **822**. As such, the current value of voltage **822** can be the current value for the reference voltage associated with the highest logical level for a TLC memory cell of the memory page. Similarly, the current value of voltage **812** can be the current value for the reference voltage associated with the second highest logical level for a TLC memory cell of the memory page.

(72) In other or similar embodiments, the reference voltage can correspond to a voltage at a particular quantile of the voltage distribution. For example, the reference voltage can correspond to a 50% quantile of the voltage distribution. As illustrated in FIG. **9A**, voltage **910** can correspond to a voltage measured at a 50% quantile of voltage distribution **810** and voltage **920** can correspond to a voltage measured at a 50% quantile of voltage distribution **820**. As such, the current value of voltage **920** can be the current value for the reference voltage associated with the highest logical level for a TLC memory cell of the memory page. Similarly, the current value of voltage **910** can be the current value for the reference voltage associated with the second highest logical level for a TLC memory cell of the memory page. It should be noted that although some embodiments and examples the present disclosure refer to processing logic (e.g., block family manager component

**113)** using a 50% quantile of a voltage distribution as a reference voltage, processing logic can use any quantile of a voltage distribution as a reference voltage. For example, processing logic can use a 25% quantile, 30% quantile 40% quantile, 75% quantile, etc. of a voltage distribution as a reference voltage.

(73) Referring back to FIG. 7, at block **730**, processing logic determines an amount of voltage shift for the memory page based on the current value for the reference voltage and a prior value for the reference voltage. The prior value for the reference voltage can be associated with a voltage distribution for the memory page at a time period before processing logic measures the current memory access voltage distribution for the memory page at block **710** (i.e., a prior voltage distribution). In some embodiments, the prior voltage distribution can correspond to a voltage distribution for the memory page immediately or soon after data for the memory page and/or the block family was programmed (i.e., written) to the memory device. In other or similar embodiments, the prior voltage distribution can correspond to a voltage distribution at any time before the processing logic measures the current voltage distribution for the memory page.

(74) FIGS. **8B** and **9B** illustrate prior voltage distributions **860**, **870** for the memory page, in accordance with embodiments of the present disclosure. In some embodiments, prior voltage distributions **860**, **870** can be specific to each memory page of each block family at the memory device. In other or similar embodiments, prior voltage distributions **860**, **870** can be specific to each block family at the memory device. For example, processing logic can associate prior voltage distributions **860**, **870** with a particular block family and, accordingly, each memory page of the block family is therefore associated with prior voltage distributions **860**, **870**. It should be noted that although some embodiments and examples of the present disclosure refer to prior voltage distributions **860**, **870** being associated with the memory page measured at block **710**, prior memory distributions **860**, **870** can be associated with sets of memory pages of any granularity (e.g., block families, voltage bins, memory device, etc.).

(75) In accordance with previously described embodiments and examples, prior voltage distributions **860**, **870** can each correspond to a voltage distribution for a distinct logical level of memory cells of the memory page. For example, voltage distribution **870** can correspond to a prior voltage distribution for a highest logical level associated with a memory cell (e.g., a TLC memory cell) and voltage distribution **860** can correspond to a prior voltage distribution for a second highest logical level associated with the memory cell.

(76) Processing logic can determine a prior value for the reference voltage in view of the prior voltage distribution **860**, **870** associated with the memory page. As described above, the reference voltage can correspond to a voltage at a particular point of a voltage distribution. For example, the reference voltage can correspond to a voltage at a tail end of a voltage distribution. In some embodiments, a voltage at a tail end of a prior voltage distribution can correspond to a voltage applied to the memory cells during a programming operations performed by the memory sub-system controller **115** to write data to the memory page (referred to as a programming voltage or a write voltage). As illustrated in FIG. **8B**, voltage **872** corresponds to the programming voltage (also referred to as a program verify (PV) voltage) used by the memory sub-system controller **115** to program data at a highest logical level at a memory cell of the memory page (e.g., PV voltage for level 7 of a TLC memory cell). Voltage **874** corresponds to a voltage associated with a tail end of voltage distribution **870** at a time period immediately or soon after programming the memory page. As seen in FIG. **8B**, the value of voltage **874** approximately corresponds to the value of voltage **872**. As the value of voltage **874** approximately corresponds to the value of voltage **872**, processing logic can use the programming voltage for the highest logical level of the memory cells of the memory page as the reference voltage. Accordingly, processing logic can determine that the prior value for the reference voltage for the highest logical level of memory cells of the memory page corresponds to voltage **872**.

(77) Voltage **862** corresponds to the programming voltage (i.e., a PV voltage) used by memory sub-

system controller **115** to program data at a second highest logical level at a memory cell of the memory page (e.g., PV voltage for level 6 of a TLC memory cell). Voltage **864** corresponds to a voltage associated with a tail end of voltage distribution **860** at the time period immediately or soon after programming the memory. As the value of voltage **864** approximately corresponds to the value of voltage **862**, processing logic can similarly use the programming voltage for the highest logical level of the memory cells of the memory page as the reference voltage. Accordingly, processing logic can determine that the prior value for the reference voltage for the second highest logical level of memory cells of the memory page corresponds to voltage **862**.

(78) By using the programming voltage as the prior value for the reference voltage for voltage distributions **860**, **870**, processing logic can determine the prior value for the reference voltage without measuring the voltage distribution of the memory pages immediately or soon after programming. For example, immediately or soon after programming, processing logic can store an indication of the voltage used to program the memory page(s) at a memory for the memory sub-system controller **115** (e.g., local memory **119**). During a calibration scan, processing logic can access the memory for the memory sub-system controller **115** to determine the programming voltage for the memory pages and use the programming voltage as the prior value for the reference voltage, in accordance with embodiments described herein.

(79) As described above, processing logic can determine an amount of voltage shift for the memory page based on the current value for the reference voltage and a prior value of the reference voltage. Referring to FIG. **8A**, processing logic can determine the amount of voltage shift **878** for voltage distribution **820** based on a difference between voltage **822** (i.e., the current value of the reference voltage associated with the highest logical memory cell level) and voltage **872** (i.e., a prior value of the reference voltage associated with the highest logical memory cell level). Similarly, processing logic can determine the amount of voltage shift **868** for voltage distribution **810** based on a difference between voltage **812** (i.e., the current value of the reference voltage associated with the second highest logical memory cell level) and voltage **862** (i.e., a prior value of the reference voltage associated with the second highest logical memory cell level).

(80) In other or similar embodiments, the reference voltage for a voltage distribution can correspond to a voltage at a particular quantile of a voltage distribution. For example, the reference voltage can correspond to a voltage at a 50% quantile of a voltage distribution. In such examples, processing logic can determine the prior value for the reference voltage based on an offset representing a difference between the voltage at a tail end of a voltage distribution and a voltage at a 50% quantile of the voltage distribution immediately or soon after programming a memory page. As described above, the voltage at a tail end of a voltage distribution can approximately correspond to a programming voltage associated with programming memory cells of a respective logical level. For example, as illustrated in FIG. **9B**, voltage **872** can correspond to a programming voltage associated with programming memory cells of a highest logical level and voltage **862** can correspond to a programming voltage associated with programming memory cells of a second highest logical level.

(81) In some embodiments, processing logic can determine the voltage at the 50% quantile of the voltage distribution based on experimental data and/or manufacturing data collected for the memory device. For example, experimental and/or manufacturing data can be collected for the memory device before the memory device is installed or initialized at the memory sub-system (e.g., at a memory device manufacturing system, a memory device testing system, etc.). The experimental and/or manufacturing data collected for the memory device can correspond to the voltage distribution for a memory page immediately or soon after programming data to each logical level at a memory cell of the memory device. Accordingly, the experimental and/or manufacturing data can indicate, for each logical level, a voltage at a tail end of a voltage distribution and a voltage at a 50% quantile of the voltage distribution immediately or soon after programming a memory page. In some embodiments, an operator of the testing system and/or a manufacturing



system can store the experimental and/or manufacturing data collected for the memory device to a memory associated with the memory sub-system controller **115** (e.g., local memory **119**). In such embodiments, processing logic can determine, for each logical level, a difference between the voltage at the tail end of the voltage distribution and the voltage at the 50% quantile of the voltage distribution (referred to as a voltage distribution offset). In other or similar embodiments, the operator can store the voltage distribution offset for each logical level at the memory associated with the memory sub-system controller **115**.

(82) Processing logic can determine the prior value for the reference voltage for a respective logical level of memory cells based on the programming voltage associated with the respective logical level and the voltage distribution offset for the respective logical level. As illustrated in FIG. **9B**, voltage **922** can correspond to a voltage at a 50% quantile of voltage distribution **870**. Processing logic can determine the value of voltage **922** based on voltage **872** (i.e., the programming voltage associated with the highest logical level of memory cells of the memory page) and the voltage distribution offset for the highest logical level of memory cells. Voltage **912** can correspond to a voltage at a 50% quantile of voltage distribution **860**. Processing logic can determine the value of voltage **912** based on voltage **862** (i.e., the programming voltage associated with the second highest logical level of memory cells of the memory page) and the voltage distribution offset for the second highest logical level of memory cells. Accordingly, processing logic can use the value of voltages **912** and **922** as the prior value for the reference voltage without measuring the voltage distribution for the memory page immediately or soon after programming at the memory device.

(83) Referring to FIG. **9A**, processing logic can determine the amount of voltage shift **924** for voltage distribution **820** based on a difference between voltage **920** (i.e., the current value of the reference voltage associated with the highest logical memory cell level) and voltage **922** (i.e., a prior value of the reference voltage associated with the highest logical memory cell level). Similarly, processing logic can determine the amount of voltage shift **914** for voltage distribution **810** based on a difference between voltage **910** (i.e., the current value of the reference voltage associated with the second highest logical memory cell level) and voltage **912** (i.e., a prior value of the reference voltage associated with the second highest logical memory cell level).

(84) Referring back to FIG. **7**, at block **740**, processing logic determines whether the amount of voltage shift for the memory page satisfies a voltage shift criterion. In some embodiments, processing logic determines that the memory page satisfies a voltage shift criterion in response to determining an amount of voltage shift determined at block **730** meets or exceeds a threshold amount of voltage shift. Similarly, processing logic can determine that the memory page does not satisfy the voltage shift criterion in response to determining the amount of voltage shift does not exceed the threshold amount of voltage shift. In response to determining that the amount of voltage shift does not satisfy the voltage shift criterion, method **700** continues to block **750**, where the processing logic associates the block family with the first voltage bin at the memory device. In response to determining that the amount of voltage shift satisfies the voltage shift criterion, method **700** continues to block **760**. At block **760**, processing logic associates the block family with a second voltage bin at the memory device that is associated with a second voltage offset. In some embodiments, processing logic can associate the block family with the second voltage bin by updating a pointer to a voltage bin boundary for the second voltage bin, as previously described. For example, processing logic can update a metadata table for the block family including the memory page to indicate that the pointer for the second voltage bin corresponds with the block family.

(85) In some embodiments, the processing logic can modify the reference voltage from a first type to a second type during operation of the memory sub-system. For example, the processing logic can use a voltage associated with the tail end of a voltage distribution as a reference voltage. During operation of the memory sub-system **110**, processing logic can monitor the performance of operations at each block family of the memory device and maintain a performance metric (e.g., a

bit error rate) for each block family. In some embodiments, processing logic can determine that the performance metric for the block family satisfies a performance criterion (e.g., falls below a performance metric threshold). In such embodiments, the processing logic can use the voltage associated with a particular quantile (e.g., a 50% quantile) of a voltage distribution as a reference voltage, in accordance with previously described embodiments.

(86) FIG. **10** is a flow diagram of another example method **1000** for voltage bin calibration based on a voltage distribution reference voltage, in accordance with some embodiments of the present disclosure. The method **1000** can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **1000** is performed by the block family manager component **113** of FIG. **1**. In other or similar embodiments, one or more operations of method **1000** is performed by another component of the memory sub-system controller, such as a block family manager. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other process flows are possible.

(87) At block **1010**, processing logic applies a programming voltage to a first set of memory cells of a memory device to program a first memory page of a block family. The processing logic applies the programming voltage to the first set of memory cells in response to determining that a memory access voltage distribution offset calibration criterion is satisfied. In some embodiments, processing logic can determine that the voltage distribution offset calibration criterion is satisfied in response to determining that an amount of time since the memory device was initialized exceeds a threshold amount of time. In other or similar embodiments, processing logic can determine that the voltage distribution offset calibration criterion is satisfied by determining that a width of a current voltage distribution for a memory page exceeds a threshold width.

(88) As described with respect to FIG. **7**, processing logic can use a voltage at a particular quantile (e.g., a 50% quantile) of a voltage distribution as a reference voltage. The prior value for the reference voltage can be determined based on a programming voltage associated with a respective logical level of memory cells at a memory page and a voltage distribution offset. The voltage distribution offset corresponds to a difference between a voltage at a tail end of a voltage distribution and the voltage at the 50% quantile of the voltage distribution at a time immediately or soon after programming of the memory cells at the respective logical level. The voltage distribution offset can be determined based on experimental and/or manufacturing data for the memory device immediately or soon after programming, the voltage distribution offset corresponds to a width of the voltage distribution for the memory device immediately or soon after programming. In some embodiments, the width of the voltage distribution for a respective logical level changes during operation of the memory device and the voltage distribution offset no longer corresponds to the distribution width.

(89) FIG. **11** illustrates a change in the width of a voltage distribution during operation of a memory device, in accordance with embodiments of the present disclosure. As described above, voltage distribution **820** represents the current voltage distribution associated with the highest logical level for memory cells of a memory page and voltage distribution **810** represents the current voltage distribution associated with the second highest logical level for memory cells of a memory page. Voltage **922** corresponds to a prior value for a reference voltage for voltage distribution **820** and voltage **912** corresponds to a prior value for a reference voltage for voltage distribution **810**. Voltage **1120** represents a current value for the reference voltage for voltage distribution **820** and voltage **1110** corresponds to a prior value for the reference voltage for voltage distribution **810**.

Voltages **1120** and **1110** are measured by processing logic in accordance with embodiments previously described.

(90) As seen in FIG. **11**, voltage **1120** does not approximately correspond to voltage **920**, which represents the current value for the reference voltage for voltage distribution **820**, as described with respect to FIG. **9A**. This indicates that a width of voltage distribution **820** is changing during operation of the memory device. Similarly, voltage **1110** does not approximately correspond to voltage **910**, which represents the current value for the reference voltage for voltage distribution **810**, as described with respect to FIG. **9**, indicating that a width of voltage distribution **810** is also increasing during operation of the memory device. Processing logic can determine that the voltage distribution offset calibration criterion is satisfied in response to determining that a difference **1112** between voltage **910** and voltage **1110** and/or a difference **1122** between voltage **920** and voltage **1120** exceed a threshold difference. A determination that the difference between voltage **910** and voltage **1110** and/or a difference between **920** and **1120** indicates that a width of a respective voltage distribution exceeds a threshold width and that the current voltage distribution offset no longer applies to voltage distributions **810** and **820**.

(91) In some embodiments, responsive to determining that the voltage distribution offset calibration criterion is satisfied, processing logic can measure the voltage distribution for the memory page and/or the block family immediately or soon after a programming operation (e.g., a write operation) is performed for the memory page and/or the block family. For example, processing logic can transmit an instruction to memory sub-system controller **115** indicating that a measurement is to be obtained for the memory page and/or block family immediately or soon after data received from the host system **120** is scheduled to be programmed to the memory page and/or the block family. The instruction can cause the memory sub-system controller **115** to measure the memory page and/or the block family in response to determining that the host data is programmed to the memory page and/or the block family. In another example, processing logic can detect that memory sub-system controller **115** has completed performance of an operation to program host data to the memory page and/or the block family and can subsequently measure the voltage distribution for memory cells of the memory page and/or the block family.

(92) At block **1020**, processing logic determines a voltage distribution offset (also referred to as a voltage distribution offset) for the block family based on the programming voltage and a first memory access voltage associated with one or more first memory cells of the first set of memory cells. The one or more first memory cells can be associated with a particular quantile (e.g., a 50% quantile) of a first voltage distribution for the memory page. Processing logic can determine the first voltage distribution for the memory page based on the measurement performed for the memory page and/or the block family, as described above. In some embodiments, processing logic can determine the voltage distribution offset by determining a difference between the programming voltage used to program the data to the memory page and/or the block family and the first memory access voltage associated with the memory cells at the particular quantile of the first voltage distribution for the memory page. In some embodiments, the first memory access voltage can correspond to the prior value for the reference voltage, in accordance with embodiments described above.

(93) At block **1030**, processing logic measures a second voltage distribution for a second memory page of the block family. Processing logic can measure the second voltage distribution for the second memory page in accordance with embodiments described with respect to FIG. **7**. For example, processing logic can measure the second voltage distribution for the second memory page at a time period after the time period at which the programming voltage was applied to program the memory cells of the memory page and/or the block family. At block **1040**, processing logic determines a second memory access voltage associated with one or more second memory cells of a second set of memory cells for the second memory page. Each of the second memory cells are associated with the particular quantile (e.g., the 50% quantile) of the second voltage distribution. In

some embodiments, the second memory access voltage can correspond to the current value for the reference voltage, in accordance with previously described embodiments.

(94) At block **1050**, processing logic determines an amount of voltage shift for the second memory page based on the second voltage and the voltage distribution offset. As described with respect to block **1020**, the voltage distribution offset corresponds to a difference between the programming voltage used to program the data to the memory page and/or the block family and the first memory access voltage associated with the memory cells at the particular quantile. Processing logic can determine the amount of voltage shift for the second memory page based on a difference between the second voltage and the voltage distribution offset, in accordance with embodiments described herein.

(95) At block **1060**, processing logic assigns the block family to a respective voltage bin of a set of voltage bins in view of the amount of voltage shift for the second memory page. As described above, processing logic can assign the block family to the respective voltage bin by associating the block family with a respective voltage bin that corresponds to the determined amount of voltage shift for the second memory page. For example, processing logic can update a pointer for a voltage bin boundary for the respective voltage bin to correspond with the block family.

(96) In some embodiments, processing logic can determine to perform the calibration scan for the block family again (e.g., in response to determining that a threshold number of P/E cycles have occurred). In such embodiments, processing logic can perform the calibration scan in accordance with embodiments described with respect to FIG. 7. During the calibration scan, processing logic can use the voltage distribution offset determined at block **1020** as the voltage distribution offset instead of the voltage distribution offset determined based on experimental and/or manufacturing data, in accordance with some embodiments.

(97) FIG. 12 illustrates an example machine of a computer system **1200** within which a set of instructions, for causing the machine to perform any one or more of the methodologies discussed herein, can be executed. In some embodiments, the computer system **1200** can correspond to a host system (e.g., the host system **120** of FIG. 1) that includes, is coupled to, or utilizes a memory sub-system (e.g., the memory sub-system **110** of FIG. 1) or can be used to perform the operations of a controller (e.g., to execute an operating system to perform operations corresponding to the block family manager component **113** of FIG. 1). In alternative embodiments, the machine can be connected (e.g., networked) to other machines in a LAN, an intranet, an extranet, and/or the Internet. The machine can operate in the capacity of a server or a client machine in client-server network environment, as a peer machine in a peer-to-peer (or distributed) network environment, or as a server or a client machine in a cloud computing infrastructure or environment.

(98) The machine can be a personal computer (PC), a tablet PC, a set-top box (STB), a Personal Digital Assistant (PDA), a cellular telephone, a web appliance, a server, a network router, a switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

(99) The example computer system **1200** includes a processing device **1202**, a main memory **1204** (e.g., read-only memory (ROM), flash memory, dynamic random access memory (DRAM) such as synchronous DRAM (SDRAM) or RDRAM, etc.), a static memory **1206** (e.g., flash memory, static random access memory (SRAM), etc.), and a data storage system **1218**, which communicate with each other via a bus **1230**.

(100) Processing device **1202** represents one or more general-purpose processing devices such as a microprocessor, a central processing unit, or the like. More particularly, the processing device can be a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor

implementing other instruction sets, or processors implementing a combination of instruction sets. Processing device **1202** can also be one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. The processing device **1202** is configured to execute instructions **1226** for performing the operations and steps discussed herein. The computer system **1200** can further include a network interface device **1208** to communicate over the network **1220**.

(101) The data storage system **1218** can include a machine-readable storage medium **1224** (also known as a computer-readable medium) on which is stored one or more sets of instructions **1226** or software embodying any one or more of the methodologies or functions described herein. The instructions **1226** can also reside, completely or at least partially, within the main memory **1204** and/or within the processing device **1202** during execution thereof by the computer system **1200**, the main memory **1204** and the processing device **1202** also constituting machine-readable storage media. The machine-readable storage medium **1224**, data storage system **1218**, and/or main memory **1204** can correspond to the memory sub-system **110** of FIG. **1**.

(102) In one embodiment, the instructions **1226** include instructions to implement functionality corresponding to a voltage bin boundary component (e.g., the block family manager component **113** of FIG. **1**). While the machine-readable storage medium **1224** is shown in an example embodiment to be a single medium, the term “machine-readable storage medium” should be taken to include a single medium or multiple media that store the one or more sets of instructions. The term “machine-readable storage medium” shall also be taken to include any medium that is capable of storing or encoding a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure. The term “machine-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media, and magnetic media.

(103) Some portions of the preceding detailed descriptions have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the ways used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

(104) It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. The present disclosure can refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage systems.

(105) The present disclosure also relates to an apparatus for performing the operations herein. This apparatus can be specially constructed for the intended purposes, or it can include a general purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program can be stored in a computer readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic-optical disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, or any type of media suitable for storing electronic instructions, each coupled to a computer system bus.

(106) The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general purpose systems can be used with programs in accordance with the teachings herein, or it can prove convenient to construct a more specialized apparatus to perform the method. The structure for a variety of these systems will appear as set forth in the description below. In addition, the present disclosure is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages can be used to implement the teachings of the disclosure as described herein.

(107) The present disclosure can be provided as a computer program product, or software, that can include a machine-readable medium having stored thereon instructions, which can be used to program a computer system (or other electronic devices) to perform a process according to the present disclosure. A machine-readable medium includes any mechanism for storing information in a form readable by a machine (e.g., a computer). In some embodiments, a machine-readable (e.g., computer-readable) medium includes a machine (e.g., a computer) readable storage medium such as a read only memory ("ROM"), random access memory ("RAM"), magnetic disk storage media, optical storage media, flash memory components, etc.

(108) In the foregoing specification, embodiments of the disclosure have been described with reference to specific example embodiments thereof. It will be evident that various modifications can be made thereto without departing from the broader spirit and scope of embodiments of the disclosure as set forth in the following claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

## Claims

1. A method comprising: determining an amount of voltage shift for one or more memory cells of a block family based on an initial reference value pertaining to the one or more memory cells and a subsequent reference value pertaining to the one or more memory cells; and associating the block family with a first voltage bin or a second voltage bin based on the determined amount of voltage shift, wherein the first voltage bin is associated with a first voltage offset and the second voltage bin is associated with a second voltage offset.
2. The method of claim 1, wherein the subsequent reference value is a current reference value and the initial reference value is associated with a memory voltage distribution for the one or more memory cells when data was written to the one or more memory cells.
3. The method of claim 1, wherein determining the amount of voltage shift for the one or more memory cells comprises: measuring a current memory access voltage distribution for the one or more memory cells, wherein the subsequent reference value is determined based on the measured current memory access voltage distribution.
4. The method of claim 3, wherein the current memory access voltage distribution corresponds to at least one of a read voltage distribution or a write voltage distribution.
5. The method of claim 1, wherein the block family is associated with the first voltage bin prior to determining the amount of voltage shift for the one or more memory cells.
6. The method of claim 5, wherein associating the block family with the first voltage bin or the second voltage bin based on the determined amount of voltage shift comprises: determining whether the determined amount of voltage shift satisfies a voltage shift criterion; and responsive to determining that the determined amount of voltage shift satisfies the voltage shift criterion, associating the block family with the second voltage bin.
7. The method of claim 1, wherein the initial reference value corresponds to a prior memory access voltage associated with the at least one of the one or more memory cells, wherein the at least one of the one or more memory cells was programmed at a faster rate than each other memory cell of the one or more memory cells.
8. The method of claim 7, wherein determining the amount of voltage shift for the one or more

memory cells comprises: obtaining a current memory access voltage for the at least one of the one or more memory cells; and determining a voltage difference between the current memory access voltage for the at least one of the one or more memory cells and the prior memory access voltage for the at least one of the one or more memory cells.

9. The method of claim 7, wherein determining the amount of voltage shift for the one or more memory cells comprises: determining a current memory access voltage associated with at least one of the one or more memory cells that is associated with a particular quantile of a current memory access voltage distribution; and determining a difference between the current memory access voltage associated with the at least one of the one or more memory cells and a prior memory access voltage associated with another memory cell associated with the particular quantile of a prior memory access voltage distribution.

10. A system comprising: a memory device; and a processing device coupled to the memory device, the processing device to perform operations comprising: determining an amount of voltage shift for one or more memory cells of a block family based on an initial reference value pertaining to the one or more memory cells and a subsequent reference value pertaining to the one or more memory cells; and associating the block family with a first voltage bin or a second voltage bin based on the determined amount of voltage shift, wherein the first voltage bin is associated with a first voltage offset and the second voltage bin is associated with a second voltage offset.

11. The system of claim 10, wherein the subsequent reference value is a current reference value and the initial reference value is associated with a memory voltage distribution for the one or more memory cells when data was written to the one or more memory cells.

12. The system of claim 10, wherein determining the amount of voltage shift for the one or more memory cells comprises: measuring a current memory access voltage distribution for the one or more memory cells, wherein the subsequent reference value is determined based on the measured current memory access voltage distribution.

13. The system of claim 12, wherein the current memory access voltage distribution corresponds to at least one of a read voltage distribution or a write voltage distribution.

14. The system of claim 10, wherein the block family is associated with the first voltage bin prior to determining the amount of voltage shift for the one or more memory cells.

15. The system of claim 14, wherein associating the block family with the first voltage bin or the second voltage bin based on the determined amount of voltage shift comprises: determining whether the determined amount of voltage shift satisfies a voltage shift criterion; and responsive to determining that the determined amount of voltage shift satisfies the voltage shift criterion, associating the block family with the second voltage bin.

16. A non-transitory computer-readable storage medium comprising instructions that, when executed by a processing device, cause the processing device to perform operations comprising: determining an amount of voltage shift for one or more memory cells of a block family based on an initial reference value pertaining to the one or more memory cells and a subsequent reference value pertaining to the one or more memory cells; and associating the block family with a first voltage bin or a second voltage bin based on the determined amount of voltage shift, wherein the first voltage bin is associated with a first voltage offset and the second voltage bin is associated with a second voltage offset.

17. The non-transitory computer-readable storage medium of claim 16, wherein the subsequent reference value is a current reference value and the initial reference value is associated with a memory voltage distribution for the one or more memory cells when data was written to the one or more memory cells.

18. The non-transitory computer-readable storage medium of claim 16, wherein determining the amount of voltage shift for the one or more memory cells comprises: measuring a current memory access voltage distribution for the one or more memory cells, wherein the subsequent reference value is determined based on the measured current memory access voltage distribution.

19. The non-transitory computer-readable storage medium of claim 18, wherein the current memory access voltage distribution corresponds to at least one of a read voltage distribution or a write voltage distribution.

20. The non-transitory computer-readable storage medium of claim 16, wherein the block family is associated with the first voltage bin prior to determining the amount of voltage shift for the one or more memory cells.

---