

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12393443
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Fishel; Liran et al.

---

### Systems and methods for task switching in neural network processor

---

#### Abstract

Embodiments relate to managing tasks that when executed by a neural processor circuit instantiates a neural network. A neural task manager circuit within the neural processor circuit can switch between tasks in different task queues. Each task queue is configured to store a reference to a task list of tasks for instantiating a neural network. Each task queue can also be assigned a priority parameter. While the neural processor circuit is executing tasks of a first task list and prior to completion of each task, the neural task manager circuit can switch between task queues according to the priority parameters for execution of tasks of a second task list by the neural processor circuit. The neural processor circuit includes one or more neural engine circuits that are configured to perform neural operations by executing the tasks assigned by the task manager.

---

<b>Inventors:</b>	<b>Fishel; Liran (Raanana, IL), Norden; Erik K. (San Jose, CA)</b>
<b>Applicant:</b>	<b>Apple Inc. (Cupertino, CA)</b>
<b>Family ID:</b>	<b>1000008762842</b>
<b>Assignee:</b>	<b>APPLE INC. (Cupertino, CA)</b>
<b>Appl. No.:</b>	<b>18/361616</b>
<b>Filed:</b>	<b>July 28, 2023</b>

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20240069957 A1	Feb. 29, 2024

#### Related U.S. Application Data

continuation parent-doc US 15971276 20180504 US 11740932 child-doc US 18361616

---

## Publication Classification

**Int. Cl.:** **G06F3/00** (20060101); **G06F9/48** (20060101); **G06F9/54** (20060101); **G06F13/16** (20060101); **G06F13/28** (20060101); **G06N3/045** (20230101); **G06N3/063** (20230101)

**U.S. Cl.:**

**CPC** **G06F9/4843** (20130101); **G06F9/544** (20130101); **G06F13/1673** (20130101); **G06F13/28** (20130101); **G06N3/045** (20230101); **G06N3/063** (20130101);

## Field of Classification Search

**CPC:** G06F (9/54)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7509671	12/2008	Bedell	718/103	G06Q 10/06
7945909	12/2010	Mariotti et al.	N/A	N/A
8407528	12/2012	Larson	N/A	N/A
9436504	12/2015	Durant	N/A	N/A
2003/0172190	12/2002	Greenblat	N/A	N/A
2006/0161696	12/2005	Anjo	710/22	G06F 15/7867
2007/0022194	12/2006	Brown	709/224	H04L 67/02
2009/0225733	12/2008	Kovacik	N/A	N/A
2009/0320034	12/2008	Tsunoda	711/E12.001	G11C 7/1006
2011/0252428	12/2010	Maruyama	718/103	G06F 9/4843
2011/0288847	12/2010	Narayanan et al.	N/A	N/A
2012/0143863	12/2011	Tran	N/A	N/A
2014/0307721	12/2013	Smith	N/A	N/A
2016/0028635	12/2015	Wang	N/A	N/A
2016/0314030	12/2015	Wilson	N/A	G06F 9/546
2016/0364644	12/2015	Brothers et al.	N/A	N/A
2016/0379109	12/2015	Chung	706/26	G06N 3/063
2017/0097625	12/2016	Waskito	N/A	G05B 19/0426
2017/0249081	12/2016	Tzori	N/A	G06F 3/0608
2017/0272304	12/2016	Castleman	N/A	N/A
2018/0188975	12/2017	Benisty	N/A	G06F 3/0679
2018/0191642	12/2017	Biederman	N/A	H04L 49/9005
2019/0057302	12/2018	Cho	N/A	G06N 3/063
2019/0102672	12/2018	Bifulco	N/A	G06N 3/04
2019/0114534	12/2018	Teng	N/A	N/A
2019/0347545	12/2018	Zhang	N/A	N/A
2019/0354496	12/2018	Nielsen	N/A	G06N 3/084
2019/0370631	12/2018	Fais et al.	N/A	N/A

2019/0392194	12/2018	Croxford et al.	N/A	N/A
2020/0019436	12/2019	Toku	N/A	G06F 9/4887
2020/0082274	12/2019	Rossi	N/A	N/A
2020/0218509	12/2019	Xu	N/A	N/A
2021/0182077	12/2020	Chen	N/A	G06N 3/045
2021/0224051	12/2020	Bequet	N/A	G06F 9/45533

## FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
106209682	12/2015	CN	N/A

## OTHER PUBLICATIONS

Lixia .Zhang, Virtual Clock: A New Traffic Control Algorithm for IPacket-Switched Networks. (Year: 1991). cited by examiner

Ranjita Bhagwan, Fast and Scalable Priority Queue Architecture for High-Speed Network Switches. (Year: 2000). cited by examiner

Anastasiya Mishchuk, Working hard to know your neighbor's margins: Local descriptor learning loss. (Year: 2017). cited by examiner

Le Nguyen Binh, A Neural-Network Contention Controller for Packet Switching Networks. (Year: 1995). cited by applicant

Mustafa K. Mehmet-Ali, The Performance Analysis and Implementation of an Input Access Scheme in a High-speed Packet Switch. (Year: 1994). cited by applicant

Uttam Kumar, A Survey on Intrusion Detection Systems for Cloud Computing Environmen, (Year: 2015). cited by applicant

Karlheinz Meier, Design and Implementation of a Multi-Class Network Architecture for Hardware Neural Networks. (Year: 2008). cited by applicant

*Primary Examiner:* Truong; Lechi

*Attorney, Agent or Firm:* Sterne, Kessler, Goldstein & Fox P.L.L.C.

## Background/Summary

(1) This application is a continuation application and claims benefit of U.S. patent application Ser. No. 15/971,276 filed on May 4, 2018, which is incorporated herein by reference in its entirety.

### BACKGROUND

#### 1. Field of the Disclosure

(1) The present disclosure relates a circuit for implementing a neural network and more specifically to managing for neural network tasks.

#### 2. Description of the Related Arts

(2) An artificial neural network (ANN) is a computing system or model that uses a collection of connected nodes to process input data. The ANN is typically organized into layers where different layers perform different types of transformation on their input. Extensions or variants of ANN such as convolution neural network (CNN), recurrent neural networks (RNN) and deep belief networks (DBN) have come to receive much attention. These computing systems or models often involve extensive computing operations including multiplication and accumulation. For example, CNN is a class of machine learning technique that primarily uses convolution between input data and kernel data, which can be decomposed into multiplication and accumulation operations.

(3) Depending on the types of input data and operations to be performed, these machine learning systems or models can be configured differently. Such varying configuration would include, for example, pre-processing operations, number of channels in input data, kernel data to be used, non-linear function to be applied to convolution result, and applying of various post processing operations. Using a central processing unit (CPU) and its main memory to instantiate and execute machine learning systems or models of various configuration is relatively easy because such systems or models can be instantiated with mere updates to code. However, relying solely on the CPU for various operations of these machine learning systems or models would consume significant bandwidth of a central processing unit (CPU) as well as increase the overall power consumption.

## SUMMARY

(4) Embodiments relate to managing tasks that when executed by a neural processor circuit instantiates a neural network. A neural task manager circuit within the neural processor circuit can switch between tasks in different task queues. Each task queue is configured to store a reference to task list of tasks for instantiating a neural network. Each task queue can also be assigned a priority parameter. As such, the neural task manager circuit can switch between task queues according to the priority parameters. The neural processor circuit also includes one or more neural engine circuits that are configured to perform neural operations by executing the tasks assigned by the task manager.

(5) Some embodiments include a method for task switching in a neural processor circuit. A reference to a first task list of first tasks that instantiates a first neural network by a neural engine circuit is stored in a first task queue circuit of a neural task manager circuit of the neural processor circuit. A reference to a second task list of second tasks that instantiates a second neural network by the neural engine circuit is stored in a second task queue circuit of the neural task manager circuit. During execution of one of the first tasks by the neural engine circuit, a task arbiter circuit, of the neural task manager circuit coupled to the first and second task queue circuits, sends configuration data for one of the second tasks, from a memory external to the neural processor circuit, for programming the neural engine circuit to instantiate at least a portion of the second neural network by executing the one of the second tasks.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Figure (FIG. 1 is a high-level diagram of an electronic device, according to one embodiment
- (2) FIG. 2 is a block diagram illustrating components in the electronic device, according to one embodiment.
- (3) FIG. 3 is a block diagram illustrating a neural processor circuit, according to one embodiment.
- (4) FIG. 4 is a block diagram of a neural engine in the neural processor circuit, according to one embodiment.
- (5) FIG. 5 is a conceptual diagram illustrating loops for processing input data at the neural processor circuit, according to one embodiment.
- (6) FIG. 6 is a conceptual diagram illustrating segmenting the input data into slices, tiles and work units, according to one embodiment.
- (7) FIG. 7 is a diagram illustrating programming of rasterizers in components of the neural processor circuit, according to one embodiment.
- (8) FIG. 8 is a flowchart illustrating a method of processing input data in a neural processor circuit, according to one embodiment.
- (9) FIG. 9 is a schematic block diagram illustrating a neural network represented by a list of tasks, according to one embodiment.

- (10) FIG. 10 is a block diagram illustrating the neural task manager, according to one embodiment.
- (11) FIG. 11 is a diagram illustrating retrieval of task descriptors using a task queue, according to one embodiment.
- (12) FIG. 12 is a diagram illustrating a task descriptor, according to one embodiment.
- (13) FIG. 13 is a block diagram illustrating a fetch queue and a configuration queue of a neural task manager, according to one embodiment.
- (14) FIG. 14 is a flowchart illustrating a method of managing tasks in a neural processor circuit, according to one embodiment.
- (15) FIG. 15 is an example neural network associated with a task queue that includes multiple tasks, according to one embodiment.
- (16) FIG. 16 is another example neural network associated with a task queue that includes multiple tasks, according to one embodiment.
- (17) FIG. 17 is a flowchart illustrating a method of task switching in a neural processor circuit, according to one embodiment.
- (18) The figures depict, and the detail description describes, various non-limiting embodiments for purposes of illustration only.

#### DETAILED DESCRIPTION

- (19) Reference will now be made in detail to embodiments, examples of which are illustrated in the accompanying drawings. In the following detailed description, numerous specific details are set forth in order to provide a thorough understanding of the various described embodiments. However, the described embodiments may be practiced without these specific details. In other instances, well-known methods, procedures, components, circuits, and networks have not been described in detail so as not to unnecessarily obscure aspects of the embodiments.
- (20) Embodiments of the present disclosure relate to managing tasks that when executed by a neural processor circuit instantiates a neural network. A machine learning operation, such as an inferencing operation or a training operation, is defined by a task list of tasks. The neural processor circuit includes one or more neural engines and a neural task manager. The neural task manager includes multiple task queues and a task arbiter. Each task queue stores a task list of tasks for a machine learning operation. Each task list or task queue may be associated with a priority parameter. The task arbiter retrieves configuration data for a task from an external memory based on the priority parameters, and provides the configuration data to components of the neural processor circuit including the one or more neural engines. In some embodiments, the neural task manager includes a configuration queue that stores configuration data of committed tasks selected by the task arbiter, and provides the configuration data to other components of the neural processor circuit. The configuration data programs the neural processor circuit to execute the task. For example, the configuration data may include input data and kernel data processed by a neural engine to execute the task. The configuration data may further include instructions for retrieving and handling the configuration data, and instructions for storing output data of the neural engine. Among other things, the neural task manager allows the neural processor circuit to efficiently handle multiple machine learning operations. Furthermore, the neural task manager may facilitate task switching when a higher priority task is stored in a task queue while a lower priority task is being executed.
- (21) Furthermore, the neural task manager, through the task arbiter, can facilitate task switching between task queues, according to the priority parameter. While the neural processor circuit is executing tasks of a first task list referenced by a first task queue and prior to execution of each of the tasks in the first task queue, the neural task manager can cause the neural processor circuit to task switch and execute one or more tasks of a second task list referenced by a second task queue. After the one or more tasks of the second task list has been executed, the neural processor circuit may return to unexecuted tasks of the first task list.
- (22) A “task” described herein refers to a processing operation of the neural processor circuit that

instantiates a network layer of a neural network, multiple network layers of a neural network, or a portion of a network layer of a neural network. A “task list” described herein refers to a sequence of tasks, such as a sequence of tasks that executed by the neural processor circuit instantiates multiple network layers of a neural network.

#### Exemplary Electronic Device

(23) Embodiments of electronic devices, user interfaces for such devices, and associated processes for using such devices are described. In some embodiments, the device is a portable communications device, such as a mobile telephone, that also contains other functions, such as personal digital assistant (PDA) and/or music player functions. Exemplary embodiments of portable multifunction devices include, without limitation, the iPhone®, iPod Touch®, Apple Watch®, and iPad® devices from Apple Inc. of Cupertino, California. Other portable electronic devices, such as wearables, laptops or tablet computers, are optionally used. In some embodiments, the device is not a portable communications device, but is a desktop computer or other computing device that is not designed for portable use. In some embodiments, the disclosed electronic device may include a touch sensitive surface (e.g., a touch screen display and/or a touch pad). An example electronic device described below in conjunction with FIG. 1 (e.g., device **100**) may include a touch-sensitive surface for receiving user input. The electronic device may also include one or more other physical user-interface devices, such as a physical keyboard, a mouse and/or a joystick.

(24) Figure (FIG. 1 is a high-level diagram of an electronic device **100**, according to one embodiment. Device **100** may include one or more physical buttons, such as a “home” or menu button **104**. Menu button **104** is, for example, used to navigate to any application in a set of applications that are executed on device **100**. In some embodiments, menu button **104** includes a fingerprint sensor that identifies a fingerprint on menu button **104**. The fingerprint sensor may be used to determine whether a finger on menu button **104** has a fingerprint that matches a fingerprint stored for unlocking device **100**. Alternatively, in some embodiments, menu button **104** is implemented as a soft key in a graphical user interface (GUI) displayed on a touch screen.

(25) In some embodiments, device **100** includes touch screen **150**, menu button **104**, push button **106** for powering the device on/off and locking the device, volume adjustment buttons **108**, Subscriber Identity Module (SIM) card slot **110**, head set jack **112**, and docking/charging external port **124**. Push button **106** may be used to turn the power on/off on the device by depressing the button and holding the button in the depressed state for a predefined time interval; to lock the device by depressing the button and releasing the button before the predefined time interval has elapsed; and/or to unlock the device or initiate an unlock process. In an alternative embodiment, device **100** also accepts verbal input for activation or deactivation of some functions through microphone **113**. The device **100** includes various components including, but not limited to, a memory (which may include one or more computer readable storage mediums), a memory controller, one or more central processing units (CPUs), a peripherals interface, an RF circuitry, an audio circuitry, speaker **111**, microphone **113**, input/output (I/O) subsystem, and other input or control devices. Device **100** may include one or more image sensors **164**, one or more proximity sensors **166**, and one or more accelerometers **168**. The device **100** may include components not shown in FIG. 1.

(26) Device **100** is only one example of an electronic device, and device **100** may have more or fewer components than listed above, some of which may be combined into a components or have a different configuration or arrangement. The various components of device **100** listed above are embodied in hardware, software, firmware or a combination thereof, including one or more signal processing and/or application specific integrated circuits (ASICs).

(27) FIG. 2 is a block diagram illustrating components in device **100**, according to one embodiment. Device **100** may perform various operations including image processing. For this and other purposes, the device **100** may include, among other components, image sensor **202**, system-on-a chip (SOC) component **204**, system memory **230**, persistent storage (e.g., flash memory) **228**,

motion (orientation) sensor **234**, and display **216**. The components as illustrated in FIG. 2 are merely illustrative. For example, device **100** may include other components (such as speaker or microphone) that are not illustrated in FIG. 2. Further, some components (such as motion sensor **234**) may be omitted from device **100**.

(28) Image sensor **202** is a component for capturing image data and may be embodied, for example, as a complementary metal-oxide-semiconductor (CMOS) active-pixel sensor, a camera, video camera, or other devices. Image sensor **202** generates raw image data that is sent to SOC component **204** for further processing. In some embodiments, the image data processed by SOC component **204** is displayed on display **216**, stored in system memory **230**, persistent storage **228** or sent to a remote computing device via network connection. The raw image data generated by image sensor **202** may be in a Bayer color kernel array (CFA) pattern (hereinafter also referred to as “Bayer pattern”).

(29) Motion sensor **234** is a component or a set of components for sensing motion of device **100**. Motion sensor **234** may generate sensor signals indicative of orientation and/or acceleration of device **100**. The sensor signals are sent to SOC component **204** for various operations such as turning on device **100** or rotating images displayed on display **216**.

(30) Display **216** is a component for displaying images as generated by SOC component **204**. Display **216** may include, for example, liquid crystal display (LCD) device or an organic light emitting diode (OLED) device. Based on data received from SOC component **204**, display **216** may display various images, such as menus, selected operating parameters, images captured by image sensor **202** and processed by SOC component **204**, and/or other information received from a user interface of device **100** (not shown).

(31) System memory **230** is a component for storing instructions for execution by SOC component **204** and for storing data processed by SOC component **204**. System memory **230** may be embodied as any type of memory including, for example, dynamic random access memory (DRAM), synchronous DRAM (SDRAM), double data rate (DDR, DDR2, DDR3, etc.) RAMBUS DRAM (RDRAM), static RAM (SRAM) or a combination thereof. In some embodiments, system memory **230** may store pixel data or other image data or statistics in various formats.

(32) Persistent storage **228** is a component for storing data in a non-volatile manner. Persistent storage **228** retains data even when power is not available. Persistent storage **228** may be embodied as read-only memory (ROM), flash memory or other non-volatile random access memory devices.

(33) SOC component **204** is embodied as one or more integrated circuit (IC) chip and performs various data processing processes. SOC component **204** may include, among other subcomponents, image signal processor (ISP) **206**, a central processor unit (CPU) **208**, a network interface **210**, sensor interface **212**, display controller **214**, neural processor circuit **218**, graphics processor (GPU) **220**, memory controller **222**, video encoder **224**, storage controller **226**, and bus **232** connecting these subcomponents. SOC component **204** may include more or fewer subcomponents than those shown in FIG. 2.

(34) ISP **206** is hardware that performs various stages of an image processing pipeline. In some embodiments, ISP **206** may receive raw image data from image sensor **202**, and process the raw image data into a form that is usable by other subcomponents of SOC component **204** or components of device **100**. ISP **206** may perform various image-manipulation operations such as image translation operations, horizontal and vertical scaling, color space conversion and/or image stabilization transformations, as described below in detail with reference to FIG. 3.

(35) CPU **208** may be embodied using any suitable instruction set architecture, and may be configured to execute instructions defined in that instruction set architecture. CPU **208** may be general-purpose or embedded processors using any of a variety of instruction set architectures (ISAs), such as the x86, PowerPC, SPARC, RISC, ARM or MIPS ISAs, or any other suitable ISA. Although a single CPU is illustrated in FIG. 2, SOC component **204** may include multiple CPUs. In multiprocessor systems, each of the CPUs may commonly, but not necessarily, implement the same

ISA.

(36) Graphics processing unit (GPU) **220** is graphics processing circuitry for performing graphical data. For example, GPU **220** may render objects to be displayed into a frame buffer (e.g., one that includes pixel data for an entire frame). GPU **220** may include one or more graphics processors that may execute graphics software to perform a part or all of the graphics operation, or hardware acceleration of certain graphics operations.

(37) Neural processor circuit **218** is a circuit that performs various machine learning operations based on computations including multiplication, addition, and accumulation. Such computations may be arranged to perform, for example, convolution of input data and kernel data. Neural processor circuit **218** is a configurable circuit that performs these operations in a fast and power-efficient manner while relieving CPU **208** of resource-intensive operations associated with neural network operations. Neural processor circuit **218** may receive the input data from sensor interface **212**, the image signal processor **206**, system memory **230** or other sources such as network interface **210** or GPU **220**. The output of neural processor circuit **218** may be provided to various components of device **100** such as the image signal processor **206**, system memory **230** or CPU **208** for various operations. The structure and operation of neural processor circuit **218** is described below in detail with reference to FIG. 3.

(38) Network interface **210** is a subcomponent that enables data to be exchanged between devices **100** and other devices via one or more networks (e.g., carrier or agent devices). For example, video or other image data may be received from other devices via network interface **210** and be stored in system memory **230** for subsequent processing (e.g., via a back-end interface to image signal processor **206**, such as discussed below in FIG. 3) and display. The networks may include, but are not limited to, Local Area Networks (LANs) (e.g., an Ethernet or corporate network) and Wide Area Networks (WANs). The image data received via network interface **210** may undergo image processing processes by ISP **206**.

(39) Sensor interface **212** is circuitry for interfacing with motion sensor **234**. Sensor interface **212** receives sensor information from motion sensor **234** and processes the sensor information to determine the orientation or movement of the device **100**.

(40) Display controller **214** is circuitry for sending image data to be displayed on display **216**. Display controller **214** receives the image data from ISP **206**, CPU **208**, graphic processor or system memory **230** and processes the image data into a format suitable for display on display **216**.

(41) Memory controller **222** is circuitry for communicating with system memory **230**.

(42) Memory controller **222** may read data from system memory **230** for processing by ISP **206**, CPU **208**, GPU **220** or other subcomponents of SOC component **204**. Memory controller **222** may also write data to system memory **230** received from various subcomponents of SOC component **204**.

(43) Video encoder **224** is hardware, software, firmware or a combination thereof for encoding video data into a format suitable for storing in persistent storage **128** or for passing the data to network interface **210** for transmission over a network to another device.

(44) In some embodiments, one or more subcomponents of SOC component **204** or some functionality of these subcomponents may be performed by software components executed on ISP **206**, CPU **208** or GPU **220**. Such software components may be stored in system memory **230**, persistent storage **228** or another device communicating with device **100** via network interface **210**.

(45) Image data or video data may flow through various data paths within SOC component **204**. In one example, raw image data may be generated from the image sensor **202** and processed by ISP **206**, and then sent to system memory **230** via bus **232** and memory controller **222**. After the image data is stored in system memory **230**, it may be accessed by video encoder **224** for encoding or by display **116** for displaying via bus **232**.

Example Neural Processor Circuit

(46) Neural processor circuit **218** is a configurable circuit that performs neural network operations



on the input data based at least on kernel data **340**. For this purpose, neural processor circuit **218** may include, among other components, neural task manager **310**, a plurality of neural engines **314A** through **314N** (hereinafter collectively referred as “neural engines **314**” and individually also referred to as “neural engine **314**”), kernel direct memory access (DMA) **324**, data buffer **318** and buffer DMA **320**. Neural processor circuit **218** may include other components not illustrated in FIG. **3**.

(47) Each of neural engines **314** performs computing operations for neural network operations in parallel. Depending on the load of operation, entire set of neural engines **314** may be operated or only a subset of the neural engines **314** may be operated while the remaining neural engines **314** are placed in a power save mode to conserve power. Each of neural engines **314** includes components for storing one or more kernels, for performing multiply-accumulate operations, and for post-processing to generate an output data **328**, as described below in detail with reference to FIG. **4**. One example of a neural network operation is a convolution operation.

(48) Neural task manager **310** manages the overall operation of neural processor circuit **218**. Neural task manager **310** may receive a task list from a compiler executed by CPU **208**, store tasks in its task queues, choose a task to perform, and send instructions to other components of the neural processor circuit **218** for performing the chosen task. Neural task manager **310** may also perform switching of tasks on detection of events such as receiving instructions from CPU **208**. In one or more embodiments, the neural task manager **310** sends rasterizer information to the components of the neural processor circuit **218** to enable each of the components to track, retrieve or process appropriate portions of the input data and kernel data, as described below in detail with reference to FIGS. **5** through **7**. Although neural task manager **310** is illustrated in FIG. **3** as part of neural processor circuit **218**, neural task manager **310** may be a component outside the neural processor circuit **218**.

(49) Kernel DMA **324** is a read circuit that fetches kernel data from a source (e.g., system memory **230**) and sends kernel data **326A** through **326N** to each of the neural engines **314**. Kernel data represents information from which kernel elements can be extracted. In one embodiment, the kernel data may be in a compressed format which is decompressed at each of neural engines **314**. Although kernel data provided to each of neural engines **314** may be the same in some instances, the kernel data provided to each of neural engines **314** is different in most instances.

(50) Data buffer **318** is a temporary storage for storing data associated with the neural network operations. In one embodiment, data buffer **318** is embodied as a memory that can be accessed by all of the neural engines **314**. Data buffer **318** may store input data **322A** through **322N** for feeding to corresponding neural engines **314A** through **314N**, as well as output from each of neural engines **314A** through **314N** for feeding back into neural engines **314** or sending to a target circuit (e.g., system memory **230**). The operations of data buffer **318** and other components of the neural processor circuit **218** are coordinated so that the input data and intermediate data stored in the data buffer **318** is reused across multiple operations at the neural engines **314**, and thereby reduce data transfer to and from system memory **230**. Data buffer **318** may be operated in a broadcast mode where input data of all input channels are fed to all neural engines **314** or in a unicast mode where input data of a subset of input channels are fed to each neural engine **314**.

(51) The input data **322** stored in data buffer **318** can be part of, among others, image data, histogram of oriented gradients (HOG) data, audio data, meta data, output data **328** of a previous cycle of the neural engine **314**, and other processed data received from other components of the SOC component **204**.

(52) Buffer DMA **320** includes a read circuit that receives a portion (e.g., tile) of the input data from a source (e.g., system memory **230**) for storing in data buffer **318**, and a write circuit that forwards data from data buffer **318** to a target (e.g., system memory).

#### Example Neural Engine Architecture

(53) FIG. **4** is a block diagram of the neural engine **314**, according to one embodiment. The neural

engine **314** performs various operations to facilitate neural network operations such as convolution, spatial pooling, and local response normalization. The neural engine **314** receives the input data **322**, performs multiply-accumulate operations (e.g., convolution operations) on the input data **322** based on stored kernel data, performs further post-processing operations on the result of the multiply-accumulate operations, and generates the output data **328**. The input data **322** and/or the output data **328** of the neural engine **314** may be of a single channel or multiple channels.

(54) Neural engine **314** may include, among other components, input buffer circuit **402**, computation core **416**, neural engine (NE) control **418**, kernel extract circuit **432**, accumulators **414** and output circuit **424**. Neural engine **314** may include further components not illustrated in FIG. 4.

(55) Input buffer circuit **402** is a circuit that stores a portion of the input data **322** as it is received from the data buffer **318** and sends an appropriate portion **408** of input data for a current task or process loop to computation core **416** for processing. Input buffer circuit **402** includes a shifter **410** that shifts read locations of input buffer circuit **402** to change the portion **408** of input data sent to computation core **416**. By changing portions of input data provided to the computation core **416** via shifting, neural engine **314** can perform multiply-accumulate for different portions of input data based on fewer number of read operations. In one or more embodiments, the input data **322** includes data of difference convolution groups and/or input channels.

(56) Kernel extract circuit **432** is a circuit that receives kernel data **326** from kernel DMA **324** and extracts kernel coefficients **422**. In one embodiment, the kernel extract circuit **432** references a look up table (LUT) and uses a mask to reconstruct a kernel from compressed kernel data **326**. The mask indicates locations in the reconstructed kernel to be padded with zero and remaining locations to be filled with numbers. The kernel coefficients **422** of the reconstructed kernel are sent to computation core **416** to populate register in multiply-add (MAD) circuits of computation core **416**. In other embodiments, the kernel extract circuit **432** receives kernel data in an uncompressed format and the kernel coefficients are determined without referencing a LUT or using a mask.

(57) Computation core **416** is a programmable circuit that performs computation operations. For this purpose, the computation core **416** may include MAD circuits MADO through MADN and a post-processor **428**. Each of MAD circuits MADO through MADN may store an input value in the portion **408** of the input data and a corresponding kernel coefficient in the kernel coefficients **422**. The input value and the corresponding kernel coefficient are multiplied in each of MAD circuits to generate a processed value **412**.

(58) Accumulator **414** is a memory circuit that receives and stores processed values **412** from MAD circuits. The processed values stored in accumulator **414** may be sent back as feedback information **419** for further multiply and add operations at MAD circuits or sent to post-processor **428** for post-processing. Accumulator **414** in combination with MAD circuits form a multiply-accumulator (MAC) **404**. In one or more embodiments, accumulator **414** may have subunits where each subunit sends data to different components of neural engine **314**. For example, during a processing cycle, data stored in a first subunit of accumulator **414** is sent to MAC circuit while data stored in a second subunit of accumulator **414** is sent to post-processor **428**.

(59) Post-processor **428** is a circuit that performs further processing of values **412** received from accumulator **414**. The post-processor **428** may perform operations including, but not limited to, applying linear functions (e.g., Rectified Linear Unit (ReLU)), normalized cross-correlation (NCC), merging the results of performing neural operations on 8-bit data into 16-bit data, and local response normalization (LRN). The result of such operations is output from the post-processor **428** as processed values **417** to output circuit **424**.

(60) NE control **418** controls operations of other components of the neural engine **314** based on the operation modes and parameters of neural processor circuit **218**. Depending on different modes of operation (e.g., group convolution mode or non-group convolution mode) or parameters (e.g., the number of input channels and the number of output channels), neural engine **314** may operate on different input data in different sequences, return different values from accumulator **414** to MAC

circuits, and perform different types of post-processing operations at post processor **428**. To configure components of the neural engine **314** to operate in a desired manner, the NE control **418** sends a control signal to components of the neural engine. NE control **418** may also include rasterizer **430** that tracks the current task or process loop being processed at neural engine **314**, as described below in detail with reference to FIG. 5 through 7.

(61) Output circuit **424** receives processed values **417** from the post-processor **428** and interfaces with data buffer **318** to store processed values **417** in data buffer **318**. For this purpose, output circuit **424** may send out as output data **328** in a sequence or a format that is different from the sequence or format in which the processed values **417** are processed in post-processor **428**.

(62) The components in the neural engine **314** may be configured during a configuration period by the NE control **418** and the neural task manager **310**. For this purpose, the neural task manager **310** sends configuration information to the neural engine **314** during the configuration period. The configurable parameters and modes may include, but are not limited to, mapping between input data elements and kernel elements, the number of input channels, the number of output channels, performing of output strides, and enabling/selection of post-processing operations at the post processor **428**.

(63) Operation of Segmenting of Data for Processing at Neural Processor Circuit

(64) Input data is typically split into smaller pieces of data for parallel processing at multiple neural engines **314**. Often multiple cycles of operations are performed to generate output for a task associated with a neural network. A compiler executed by CPU **208** analyzes the hierarchy and nodes of the neural network and determines how the input data is to be segmented based on the hardware constraints of the neural processor circuit **218**. One of functions of the compiler is to determine how input data is to be split into smaller data units for processing at the neural engines **314**, and how the processing is to be iterated in loops to produce the result for tasks.

(65) FIG. 5 is a conceptual diagram illustrating loops for processing the input data at neural processor circuit **218**, according to one embodiment. The outermost loop represents processing for a convolution group, if group convolution involving multiple convolution group is used. Group convolutions are convolutions where input data of the input channels in each group are used only for generating output data of output channels of each group but are not used for generating output data for output channels of other groups. Hence, each group of the group convolution can be treated as a separate convolution operation.

(66) In the loop for each convolution group is a processing loop for a slice of the input data. The entire input data for a convolution operation is segmented into multiple strips of slices in an overlapping manner, as shown in FIG. 6. The overlapping portions **602**, **604**, **606** are parts of the input data that are overfetched in two adjacent slices to provide spatial support for a corresponding kernel. The second outermost loop performs convolution operation for each slice in the input data. Within the loop for a slice is a processing loop for a tile of the slice. Each slice is segmented into a plurality of tiles, as shown in FIG. 6. The overlapping portions **608**, **610**, **612**, **614** are parts of the input data in slice **4** that are overfetched in two adjacent tiles to provide spatial support for a corresponding kernel. The rightmost tile will typically have a width smaller than other tiles of the slice. In one embodiment, input data for each tile is loaded onto data buffer **318** in a read cycle and reused for operations in processing loops for the tile. In the processing loop for the tile is a processing loop for a work unit. Each tile is segmented into multiple work units as shown in FIG. 6. A work unit is a portion of the input data having a size that produces output values that fit into accumulator **414** of neural engine **314** during a single cycle of the computation core **416**. Although the shape of each work unit is shown as a horizontal strip in FIG. 6, the shape of the work unit can be different depending on the shape and size of the tile. The work units also have overlapping parts that represent overfetched data to provide support for a corresponding kernel. Especially, work units for the last tile of a slice may have a shape of a vertical strip if the tile is tall. In one or more embodiments, the size of each work unit is 256 bytes. In such embodiments, for example, work

units be shaped to one of 16×16, 32×8, 64×4, 128×2 or 256×1 dimension.

(67) For each work unit, an internal processing loop may be provided for an output channel group (OCG). The number of output channels produced for a given work unit by a single cycle of the computation core **416** is referred to as an OCG. Depending on operation modes, each neural engine **314** may process output data of different numbers of output channels (e.g., 8 channels, 32 channels) for a single load of input data into its input buffer circuit **402**.

(68) For each output channel group, an internal processing loop may be provided for an input channel (Cin). If an input stride is implemented to skip certain input data, loops for sub-input channels (Sub-Cin) may be provided within the processing loop for the input channel (Cin).

(69) For each input channel or each sub-input channel, internal loops are provided for processing horizontal spatial support for a kernel and the vertical support within each horizontal spatial support. The spatial support refers to the input data for convolution with the kernel, and includes overfetched input data for performing convolution at the edges of the input data.

(70) Overfetch refers to fetching additional input data in current slice, tile or work unit so that proper dimension of input data can be provided for convolution with a kernel. In one or more embodiments, overfetch is performed vertically between slices to obtain additional rows of input data (shown as overlapping portions **602**, **604**, **606** in FIG. 6), horizontally between tiles to obtain additional columns of input data (shown as overlapping portions **608**, **610**, **612**, **614** in FIG. 6), and vertically between work units within a tile to obtain additional rows of input data.

(71) For each spatial support for the kernel, an internal processing loop for an output channel (OC) is provided to generate output data for each output channel (Cout). In cases where output stride implements a spatial upsampling, an additional inner loop for processing each sub-output channel is provided. Loading of kernel coefficients and MAC operations are performed within the loop for the output channel (OC) or sub-output channel if an output stride is implemented, to generate output data for the output channel (OC) or sub-output channel.

(72) The nested loop structure of FIG. 5 is merely illustrative. Loops may be omitted, added or structured differently depending on various factors. For example, if only a single convolution group is used, the outermost loop may be removed. Further, the loop structure for the horizontal spatial support and the vertical spatial support may be reversed.

(73) In one or more embodiments, the operations associated dividing the input space into smaller units and processing these smaller units as described above with reference to FIGS. 5 and 6 are performed by rasterizers **714**, **718**, **720**, **722** in various components of neural processor circuit **218**. A rasterizer is a circuit in various components of neural processor circuit **218** that keeps track of the segment of the input/output data (e.g., group, work unit, input channel, output channel) and instructs the components of neural processor circuit for proper handling of the segment of the input data. For example, rasterizer **720** in buffer DMA **320** tracks tiles and slices received from system memory **230** while rasterizer **718** in data buffer **318** broadcasts in sequence work units for processing by the neural engines **314**. Rasterizer **722** in kernel DMA **324** determines which kernels are to be received and distributed to neural engines **314**, while rasterizers **714** in neural engines **314** operate shifters **410** in input buffer circuits **402** to forward correct portions **408** of input data to MAC **404**, and send the finished output data **328** to the data buffer **318**.

(74) FIG. 7 is a diagram illustrating programming of rasterizers **714**, **718**, **720**, **722** in components **314**, **318**, **320**, **322** of the neural processor circuit **218**, according to one embodiment. To perform their functions, each of rasterizers **714**, **718**, **720**, **722** receives task information **710** (e.g., configuration data) indicating how the input data and/or kernel data are to be segmented and to be handled by each component of the neural processor circuit **218**. The task information includes information about particulars of the current layer (e.g., dimensions of input and output data, dimension of an associated kernel, types of padding at the boundaries of input data). Rasterizers **714**, **718**, **720**, and **722** may also receive constraints on their operations (e.g., whether to allow or disallow tile width over a threshold).

(75) By providing rasterizers in different components of neural processor circuit **218**, overhead in data transmitted between the components of the neural processor circuit **218** may be reduced. If a single central rasterizer is provided to control different components of the neural processor circuit **218**, kernel data, input data, and output data transmitted between the components may be needed in these data to identify associated position in the loops of the task such as convolution group, tile, slice, work unit, input channel and output channel. By using distributed rasterizers, no separate metadata is needed to transmit the kernel data, input data and output data among components of the neural processor circuit **218**.

#### Example Process at Neural Engine Architecture

(76) FIG. **8** is a flowchart illustrating a method of processing input data in neural processor circuit **218**, according to one embodiment. The method may include different and/or additional steps, or the steps may be in different orders.

(77) After neural task manager **310** programs rasterizers **714**, **718**, **720**, **722**, the process of operating buffer DMA **320** is initiated by rasterizer **720** instructing **804** buffer DMA **320** to cause buffer DMA **320** to receive a tile of input data from system memory **230**. The tile received by buffer DMA **320** is stored **806** in data buffer **318**.

(78) Rasterizer **718** in data buffer **318** then instructs **808** data buffer **318** to send a work unit to one or more neural engines **314**. The work unit is then stored in input buffer circuits **402** of the one or more neural engines **314**.

(79) In one or more embodiments, input buffer circuit **402** selects **816** a portion of work unit to be sent to MAC **404** to perform multiply-accumulate operation. Then MAC **404** performs **820** multiply-accumulate operations on the selected portion of the work unit using a corresponding kernel. Then it is determined **824** if the entire work unit is processed at one or more neural engines **314**. If not, the selected portion of the work unit is shifted **828** by shifter **410** and returns to perform **820** another round of multiply-accumulate operations.

(80) If it is determined **824** that the entire work unit was processed, then it proceeds to determine **832** if all work units in the tile was processed. If not, then the process proceeds **836** to the next work unit by having data buffer **318** send **808** a next work unit to one or more neural engines **314**, and repeats the subsequent processes.

(81) If it is determined **832** that all work units in the tile was processed by the neural engines **314**, the process proceeds to determine **840** whether all tiles for the input data were processed. If not, the process proceeds **844** to a next tile by having rasterizer **720** instructs **804** buffer DMA **320** to receive a next tile from system memory **230** and repeats the subsequent processes.

(82) If it is determined **840** that all tiles of the input data are processed, then the process ends for the current input data. Then, the process may be repeated to process the next input data or proceed to the next task.

(83) Embodiments of the process as described above with reference to FIG. **8** are merely illustrative. Further loops may be embodied, as described above with reference to FIG. **5**. Moreover, sequence of the process may be modified or omitted.

#### Example Neural Task Manager Architecture

(84) A neural network may include network layers or sub-layers that are instantiated or implemented as a series of tasks executed by the neural processor circuit **218**. FIG. **9** is a schematic block diagram illustrating a neural network **900** represented by a list **904** of tasks, according to one embodiment. The neural network **900** includes network layers (or sub-layers) including convolution layers C1, C2, C3 (including sub-layers C300, C310, C311, C320, and C321), C5, and C6, and pooling layers P2 and P4. The neural network **900** is an example of a neural network architecture that may be instantiated by the neural processor circuit **218**. That is, when the neural network **900** is converted into the task list **904** to become executable by the neural processor circuit **218**. Other types of neural network architectures with different types of network layers or orders of network layers may also be instantiated by the neural processor circuit **218**.

(85) The neural network **900** is converted, such as by the CPU **208**, to the task list **904**. The task list **904** includes a linear link-list defining a sequence of tasks including task C1, task C2+P2, task C3.sub.00+P4, task C3.sub.10, task C3.sub.11+P4, task C3.sub.20, task C3.sub.21+P4, task C5.sub.a, task C5.sub.b, and task C6. Each task is associated with a task descriptor that defines a configuration of the neural processor circuit **218** to execute the task. Each task may correspond with a single network layer of the neural network **900**, a portion of a network layer of the neural network **900**, or multiple network layers of the neural network **900**. For example, the task C1 corresponds with the single network layer C1, the task C2+P2 corresponds with multiple network layers C2 and P2, and the tasks C5.sub.a and C5.sub.b each correspond with a portion of the network layer C5. The neural processor circuit **218** instantiates the neural network **900** by executing the tasks of the task list **904** under the control of the neural task manager **310**.

(86) FIG. **10** is a block diagram illustrating the neural task manager **310**, according to one embodiment. The neural task manager **310** manages the execution of tasks for one or more neural networks **900** by the neural processor circuit **218**. The neural task manager **310** may include, among other components, a task arbiter **1002**, task queues **1004A** through **1004N** (hereinafter collectively referred to as “task queues **1004**” and individually also referred to as “task queue **1004**”), a task manager direct memory access (DMA) **1006**, a fetch queue **1008**, and a configuration queue **1010**. The neural task manager **310** may include other components not illustrated in FIG. **10**.

(87) The task arbiter **1002** is a circuit or a combination of circuit and firmware that selects tasks from the task queues **1004** for execution by the neural processor circuit **218**. The task arbiter **1002** dequeues tasks from the task queues **1004**, and places tasks in the configuration queue **1010**. While a task is in a configuration queue, it is committed to execution and the neural processor circuit performs a prefetch for input data and kernel data before the task is executed by other components of the neural processor circuit **218**. For example, the task arbiter **1002** may perform fixed-priority arbitration between multiple task queues **1004**, and select the task from the task queues **1004** with the highest priority for retrieval of a task descriptor **1012** from the system memory **230** by the task manager DMA **1006**.

(88) The neural task manager **310** may include one or more task queues **1004**. Each task queue **1004** is coupled to the CPU **208** and the task arbiter **1002**. Each task queue **1004** receives from the CPU **208** a reference to a task list **904** of tasks that when executed by the neural processor circuit **218** instantiates a neural network **900**. The reference stored in each task queue **1004** may include a set of pointers and counters pointing to the task list **904** of task descriptors **1012** in the system memory **230**. Each task queue **1004** may be further associated with a priority parameter that defines the relative priority of the task queues **1004**. The task descriptor **1012** of a task specifies a configuration of the neural processor circuit **218** for executing the task.

(89) The task manager DMA **1006** is coupled to task arbiter **1002**, the system memory **230**, and the fetch queue **1008**. The task manager DMA **1006** includes a read circuit that receives task descriptors **1012** of tasks from a source (e.g., system memory **230**) for storing in the fetch queue **1008**. For example, the task arbiter **1002** selects a task queue **1004** according to the priorities of the task queues **1004**, and uses the task list **904** referenced by the selected task queue **1004** to control the task manager DMA **1006** to select the task descriptor **1012** of a task.

(90) The fetch queue **1008** is a single entry queue that stores a task descriptor **1012** of a task that is pending to commit for execution. The fetch queue **1008** is coupled to the task manager DMA **1006** to receive the task descriptor **1012** from the system memory **230**, and provides the task descriptor **1012** to the configuration queue **1010** or the configuration data **1014** extracted from the task descriptor **1012** to the configuration queue **1010**.

(91) The configuration queue **1010** holds configuration data **1014** of multiple tasks that have been committed for execution. When a task is in the configuration queue **1010**, the kernel DMA **324** may fetch kernel data from the system memory **230** to store in the kernel extract circuit **432** of neural engines **314**, and the buffer DMA **320** may fetch input data from the system memory **230** to store in

the data buffer **318**. To execute the task, the kernel extract circuit **432** provides the prefetched kernel data to the MAC **404** of the neural engine **314**, and the data buffer **318** provides the prefetched input data to the MAC **404** of the neural engine **314**. In some embodiments, the configuration queue **1010** may include multiple queues that holds configuration data **1014** extracted from the committed task descriptors **1012**. As discussed in greater detail in connection with FIG. **13**, the configuration queue **1010** is further coupled to other components of the neural processor circuit **218** to configure the neural processor circuit **218** according to the configuration data **1014**.

(92) FIG. **11** is a diagram illustrating retrieval of task descriptors **1012** using a task queue **1004**, according to one embodiment. The task queue **1004** includes a reference, such as a set of pointers, to the task descriptors **1012A** through **1012N** stored in the system memory **230**. To that end, the task queue **1004** may include a memory storing a head parameter **1102**, a network identifier (ID) **1104**, a base address index **1106**, a tail parameter **1108**, a count parameter **1110**, a priority parameter **1112**, and a task queue switch parameter **1116**. The head parameter **1102** is a pointer to a location of the system memory **230** storing the task descriptor **1012A** at the head of the task queue **1004**. The network ID **1104** identifies the neural network **900** of the task descriptor **1012** at the head of the task queue **1004**, and the base address index **1106** is an index to a base-address table **1114** tagged with the network ID **1104** of the task descriptor **1012A** at the head of the task queue **1004**. The count parameter **1110** defines the number of task descriptors **1012** in the task queue **1004**. The priority parameter **1112** defines the priority of the task queue **1004**, which is used by the task arbiter **1002** to select between multiple task queues **1004**. The task queue switch parameter **1116** indicates whether the task queue **1004** was interrupted due to a task switch and needs to be completed.

(93) When a particular task queue **1004** is selected (e.g., according to the priority parameter **1112**), the task arbiter **1002** references the head parameter **1102**, the network ID **1104**, the base address index **1106**, and the base address table **1114** to retrieve a task descriptor **1012** from the system memory **230**, and places the task descriptor **1012** into the fetch queue **1008** to initiate commitment of the task for execution. In each configuration period, the task arbiter **1002** may continue to place a task descriptor **1012** into the fetch queue **1008** according to the order of tasks defined by the task list **904** of the task queue **1004**, such as by retrieving the next task descriptor **1012B**, and so forth.

(94) FIG. **12** is a diagram illustrating a task descriptor **1012**, according to one embodiment. The task arbiter **1002** places the task descriptor **1012** in the fetch queue **1008** from system memory **230**, which is then transferred to the configuration queue **1010**. The highest priority (e.g., first in) task descriptor **1012** in the configuration queue **1010** is used to configure the neural processor circuit **218** for execution during the configuration period. The task descriptor **1012** includes configuration data **1014** including a task descriptor header **1202** and address data **1204A** through **1204N** (hereinafter referred to as “address data **1204**”). The task descriptor header **1202** includes configuration data **1014** that configures various operations of the neural task manager **310**, including operations related to task selection and task switching. For example, the task descriptor header **1202** may be parsed by the task arbiter **1002** to extract configuration data **1014** that programs the neural task manager **310** and other components of the neural processing circuit **218**. The task descriptor header **1202** may include a task identifier (ID) **1206** that identifies the task, a neural network identifier (ID) **1208** that identifies a neural network **900** instantiated by the task, a task switch parameter **1210** defining whether the neural task manager **310** should initiate a task switch (e.g., to execute a task of a different task queue **1004**) after execution of the task, an input surface parameter **1212** defining whether the input data for the task should be retrieved from the system memory **230** or the data buffer **318**, an output surface parameter **1214** defining whether the output data of the task should be stored in the system memory **230** or the data buffer **318**, various (e.g., base address) pointers **1216** to facilitate the programming of the neural processor circuit **218**, and one or more debug/exception parameters **1222** that control event, exception, or debug logging. The task descriptor header **1202** may also include additional parameters related to task switching,

such as a task switch ready (TSR) parameter **1218** that defines whether the neural task manager **310** can task switch after execution of the task, and a source pointer last (SPL) parameter **1220** that indicates, after returning to an interrupted task queue **1004**, the task is the last task with input data stored in the system memory **230**. Parameters related to task switching will be further described with reference to task switching (e.g., see descriptions related to FIGS. **15-17**).

(95) Each instance of address data **1204A** through **1204N** (collectively or individually referred to as “address data **1204**”) defines an address and data payload pair used to program the components of the neural processor circuit **218**. The data payload may include input data and kernel data used to execute the task. For example, each instance of address data **1204** includes register data defining the data payload, a register address defining a destination memory location of the neural processing circuit **218** for receiving the register data, and a register count defining a number of consecutive memory locations (e.g., registers) to be written with the register data. In some embodiments, the register address is combined with the base address stored in the header **1202** to define the full address of each memory location. If the task descriptor **1116** is generated at compile time, then the actual run time addresses may not be known. The base address table **1114** is used avoid duplicating or updating all task descriptors with dynamically assigned addresses.

(96) FIG. **13** is a block diagram illustrating the fetch queue **1008** and configuration queue **1010**, according to one embodiment. The configuration queue **1010** is coupled to the fetch queue **1008**, which is coupled to the system memory **230** via the task manager DMA **1006**. The configuration queue **1010** is further coupled to the rasterizer **714** of one or more neural engines **314**, the rasterizer **718** of the data buffer **318**, the rasterizer **720** of the buffer DMA **320**, and the rasterizer **722** of the kernel DMA **324**. The fetch queue **1008** stores a task descriptor **1012** (e.g., including the task descriptor header **1202** and the address data **1204A** through **1204N**) for a task that is pending and not committed to execution. The fetch queue **1008** reduces the latency of reading the next task descriptor **1012** into the configuration queue **1010** from the system memory **230**. The fetch queue **1008** stores the highest priority task descriptor **1012** as determined by the task arbiter **1002**. The task arbiter **1002** may replace the task descriptor **1012** stored in the fetch queue **1008** if a higher priority task descriptor **1012** has been has been enqueued (e.g., from a higher priority task queue **1004**). The task descriptor **1012** in the fetch queue **1008** does not initiate an input data or kernel prefetch, and does not affect task queue priorities, pointers, or counters. As such, a task descriptor **1012** in the fetch queue **1008** may be readily replaced by a higher priority task descriptor **1012** by writing the higher priority task descriptor **1012** into the fetch queue **1008**. When a task descriptor **1012** stored in the configuration queue **1010** is executed by the neural processor circuit **218**, the task descriptor **1012** stored in the fetch queue **1008** is transferred to the configuration queue **1010**, and another task descriptor **1012** of a subsequent task may be stored in the fetch queue **1008**.

(97) The configuration queue **1010** stores task descriptors **1012** of tasks committed for execution by the neural processor circuit **218**. In some embodiments, the configuration queue **1010** includes multiple separate queues **1310** that each store a portion of the configuration data **1014** (including configuration data **1014A** through **1014E**) extracted from the task descriptor **1012**. Furthermore, the queues **1310** are each coupled to a respective component of the neural processor circuit **218** for programming the component with the configuration data **1014**. Through operation of the configuration queue **1010**, the neural task manager **310** programs the rasterizers **714**, **718**, **720**, **722** to perform the functionality discussed above in FIGS. **7** and **8**. For example, a queue **1310A** is coupled to the rasterizers **714** of the neural engines **314** to provide configuration data **1014A** that controls the operations of the shifters **410** in the input buffer circuits **402** to forward correct portions **408** of input data to MAC **404**, and send the finished output data **328** to the data buffer **318**. The queue **1310B** is coupled to the rasterizer **718** of the data buffer **318** to provide configuration data **1014B** that controls the broadcasting of input data (e.g., work units) by the data buffer **318** for processing by the neural engines **314**. The queue **1310C** is a read queue that is coupled to the rasterizer **720** of the buffer DMA **320** to provide configuration data **1014C** that



controls the buffer DMA **320** to retrieve input data (e.g., a tile) from system memory **230** and store the input data in the data buffer **318**. The queue **1310D** is a write queue that is coupled to the rasterizer **720** of the buffer DMA **320** to provide configuration data **1014D** that controls the buffer DMA **320** to store output data in the system memory **230**. The queue **1310E** is coupled to the rasterizer **722** of the kernel DMA **322** to provide configuration data **1014E** that controls which kernels are to be received and distributed to neural engines **314**. In some embodiments, a task descriptor **1012** or configuration data **1014** stored in the configuration queue **1010** cannot be replaced and will be executed in a first in, first out (FIFO) order.

#### Example Process at Neural Task Manager Architecture

(98) FIG. **14** is a flowchart illustrating a method of managing tasks in the neural processor circuit **218**, according to one embodiment. The method may include different and/or additional steps, or the steps may be in different orders.

(99) The CPU **208** generates **1402** a task list **904** of task descriptors **1012** of tasks that when executed by the neural processor circuit **218** instantiates a neural network **900**. For example, the CPU **208** may receive input data for a machine learning operation from the image sensor **202**, the system memory **230**, the persistent storage **228**, the network interface **210**, or some other components. The machine learning operation may include an inferencing operation, or a training operation. The neural network **900** may include kernel data and a neural network architecture including network layers. The input data is applied to the neural network **900** to perform the machine learning operation. The kernel data and network layers of the neural network **900** may be computed in a machine learning training process. The CPU **208** performs a compile operation (offline or on-the-fly) to turn a neural network description **900** into a linked list of task descriptors **1012** referred to herein as the task list **904**. Each task is defined by a task descriptor **1012**, and when executed by the neural processor circuit **218** instantiates a single network layer, multiple network layers, or a portion of a network layer. Each task descriptor **1012** of a task includes configuration data **1014**, such as the task descriptor header **1202** and the address data **1204** defining neural processor circuit **218** address and data payload pairs. The data payload may include the kernel data of the neural network **900** and the input data. The configuration data **1014** further includes instructions that configure operation of the rasterizers **714**, **718**, **720**, and **722** to execute the task.

(100) The CPU **208** stores **1404** the task list **904** of task descriptors **1012** in the system memory **230**. In some embodiments, the CPU **208** or another CPU external to the electronic device **100** generates the task descriptors **1012** and stores the task descriptors **1012** in the persistent storage **228** or some other non-volatile memory. The task descriptors **1012** are loaded along with the kernel data and input data in the system memory **230** for use by the neural processor circuit **218**. The CPU **208** may be coupled to the system memory **230** via the bus **232** and the memory controller **222**.

(101) The task list **904** may include a set of pointers that reference the locations of the task descriptors **1012** in the system memory **230**. Furthermore, the CPU **208** may update parameters of the task descriptors **1012** such as memory addresses or the network ID. For example, the task list **904** may include a head parameter **1102** and a tail parameter **1108** respectively defining a beginning register address and an end register addresses of the system memory **230** where multiple tasks descriptors **1012** are stored. In some embodiments, the references to the register addresses in the task list **904** may be partial addresses, and a base address table **1114** is used to define the full reference addresses to the system memory **230**. In some embodiments, the CPU **208** may patch absolute addresses as necessary. The CPU **208** may further configure the neural processor circuit **218** by setting its base address registers.

(102) The task arbiter **1002** of the neural task manager **310** enqueues **1406** the task list **904** for execution by the neural processor circuit **218**. For example, the neural task manager **310** includes multiple task queues **1004A** through **1004N**. The task queues **1004** may each store a reference to a task list **904**. Furthermore, the task queues **1004** are prioritized for execution according to the

priority parameters **1112** of the task lists **904** referenced by the task queues **1004**.

(103) In some embodiments, the CPU **208** performs a general configuration of the neural processor circuit **218** to execute the task. The CPU **208** may further start the neural processor circuit **218** if the neural processor circuit **218** is not already running.

(104) The task arbiter **1002** selects **1408** a task queue with highest priority for execution. For example, the task arbiter **1002** in each programming period selects a task queue **1004** based a comparison of priority parameters **1112** of the task queues **1004** or the task lists **904** of the task queues **1004**, and executes tasks from the task list **904** from the highest priority task queue **1004**.

(105) The neural processor circuit **218** executes **1410** a task from the selected task queue. For example, the neural processor circuit **218** performs the method of FIG. **8** to execute the task. The configuration queue **1010** of the neural task manager **310** may provide the configuration data **1014** (or task information **710**) to the rasterizers **714**, **718**, **720**, **722** of the neural processor circuit **218** to program the neural processor circuit **218** to execute the task, as shown in FIG. **7**. Furthermore, the execution of the task may include processing of the prefetched kernel data and input data while the task was in the configuration queue **1010**. The execution of each task may include multiple processing loops for handling the input data as shown in FIGS. **5** and **6**. The task arbiter **1002** may dequeue the task from the configuration queue **1010** after execution of the task.

(106) In some embodiments, to execute the task, the task arbiter **1002** controls the task manager DMA **1006** to retrieve the task descriptor **1012** of the task of the task list **904**, and store the task descriptor **1012** in the fetch queue **1008**. After execution of a previously committed task, the task arbiter **1002** may dequeue the executed task by removing the task descriptor or configuration data of the task from the configuration queue **1010**. The task descriptor **1012** or extracted configuration data **1014** of the current task is then placed in the configuration queue **1010** from the fetch queue **1008**. When the task is in the configuration queue **1010**, the neural processor circuit **218** may initiate a prefetch operation by the kernel DMA **324** for kernel data from the system memory **230** to the kernel extract circuit **432**, and a prefetch operation by the buffer DMA **320** for input data from the system memory **230** to the data buffer **318**.

(107) If a different task queue has a higher priority, a task switch to the task queue with the higher priority may be performed when possible. To perform a task switch (e.g., to another machine learning operation), the task arbiter **1002** replaces a task in the fetch queue **1008** with another task descriptor referenced in a different task queue **1004**. The task arbiter **1002** may resume the interrupted machine learning operation defined by the task list **904** after completion of the task switch by storing the replaced task into the fetch queue **1008**.

(108) In some embodiments, the CPU **208** determines **1408** a priority parameter **1110** of the task list **904** to select a task queue. The priority parameter **1110** defines the priority of the tasks of the machine learning operation relative to other tasks of other machine learning operations executed by the neural processor circuit **218**. The CPU **208** may facilitate the execution of multiple machine learning operations. The machine learning operations may be different, such as by using different input data or different neural network architectures and kernel data. The CPU **208** may determine the priority parameter of the task list **904** programmatically, or based on user input. The task descriptor **1012** referenced by the task queue **1004** with the highest priority may be selected for execution. The priority of a task queue is either determined by the CPU **208**, or dynamically by information from a previously executed task descriptor.

(109) The task arbiter **1002** determines **1412** whether all tasks of the task list **904** have been executed. If a task of the task list **904** has not been executed, the process returns to step **1410**, where the task arbiter **1002** executes the unexecuted task of the task list **904**.

(110) If each task of the task list **904** has been executed, the task arbiter **1002** removes **1414** the task list **904** from the task queue **1004**. In another example, a reference to a task descriptor **1012** of a task may be removed from the task queue **1004** subsequent to execution of the task. The process may end. In another example, the process may return to **1402**, where the CPU **208** may continue to

generate task descriptors of tasks for other machine learning operations. In that sense, the process may be repeated to place tasks into task queues for execution of machine learning operations by the neural processor circuit **218** according to specified task priorities.

#### Example Process at Neural Task Manager Task Switching

(111) The neural task manager **310**, through the task arbiter **1002**, can support task switching between task queues **1004**. While the neural processor circuit **218** is executing tasks of a first task list of a first task queue **1004** and prior to execution of each of the tasks in the first task queue **1004**, the neural task manager **310** can cause the neural processor circuit **218** to task switch and execute one or more tasks of a second task list in a second task queue **1004**. After the one or more tasks of the second task list has been executed, the neural processor circuit **218** may return to unexecuted tasks of the first task list. For example, the neural task manager **310** provides configuration data for a task of the second task list to one or more neural engines **314** for execution. Task switching may occur due to a task queue **1004** having a higher priority parameter **1112** than a task queue **1004** that is being executed. The priority parameter **112** of a task queue **1004** can be determined according to the task urgency, task length, task source, etc. For example, to execute a high priority machine learning operation, the CPU **208** may generate a task list and store a reference to the task list in a task queue **1004** with a higher priority parameter **1112** than a task list of a task queue **1004** that is being executed.

(112) Without a task switch, output data **328** of a task in a task queue **1004** may be stored in the data buffer **318** (e.g., a destination target is in the data buffer **318**) and may be used as input data for a subsequent task in the task queue **1004** (e.g., a source target for the subsequent task is in the data buffer **318**). However, during a task switch, source targets (e.g., for input data) and destination targets (e.g., for output data **328**) may change from the data buffer **318** to the system memory **230**. In some embodiments, data is stored in a memory cache (e.g., instead of the system memory **230**) between the data buffer **328** and system memory **230** (not shown in FIG. 2). For example, during a task switch from a first task queue **1004** being executed to a second task queue **1004**, output data **328** for executed tasks in the first task queue **1004** are stored in the system memory **230**, instead of the data buffer **318**. In another example, after each of the tasks of the second task queue **1004** are executed, the output data **328** of the tasks of the second task queue **1004** are stored in the system memory **230** and the input data for the remaining tasks in the first task queue **1004** are retrieved from the system memory **230**.

(113) As previously described with reference to FIG. 12, task descriptor headers **1202** can include configuration data **1014** that configures behaviors related to task switching. Specifically, the task switch enable parameter **1210** (TSE) defines whether the neural task manager **310** can begin a task switch process after execution of the task (e.g., a task switch process can begin if TSE=1). The input surface parameter **1212**, also referred to as the source pointer change parameter (SPC), defines whether the input data for the task should be retrieved from the system memory **230** or the data buffer **318** (e.g., input data is retrieved from the system memory **230** if SPC=1). The output surface parameter **1214**, also referred to as a destination pointer change parameter (DPC) defines whether the output data **328** of the task should be stored in the system memory **230** or the data buffer **318** (e.g., the output data **328** is stored in the system memory **230** if DPC=1).

(114) Task descriptor headers **1202** can also include a task switch ready parameter **1218** (TSR) that defines whether the neural task manager **310** can task switch after execution of the task (e.g., task switch if TSR=1) and a source pointer last parameter **1220** (SPL) that indicates, after returning to an interrupted task queue **1004**, the task is the last task with input data stored in the system memory **230**. Furthermore, each task queue **1004** can have a task queue switch parameter **1116** (TQSP) that indicates whether the task queue **1004** was interrupted due to a task switch and needs to be completed (e.g., TQSP=1 indicates that the task queue **1004** was interrupted due to a task switch). The neural task manager **310** may also have a global task switch parameter (GTSP) that indicates a task switch is in process (e.g., GTSP=1 indicates a task switch is in progress).

(115) The neural task manager **310** can follow a set of rules that configures the neural task manager **310** to direct a task switch process. These rules can include:

(116) Rule 1: if a task switch request is received and TSE=1 for the task, then GTSP is set to 1, according to an embodiment. A task switch request can be triggered by the task manager **310** after determining that a task or a task queue **1004** has a higher priority parameter **1112** than a current task or current task queue **1004**.

(117) Rule 2: if GTSP=1 and DPC=1, then the output data **328** from the executed task is stored to the system memory **230**, according to an embodiment.

(118) Rule 3: if GTSP=1 and TSR=1, then TQSP is set to 1 for the current task queue **1004** (i.e., the task queue being executed), GTSP is set to 0, and the task switch can occur, according to an embodiment.

(119) Rule 4: if SPC=1 and TQSP=1, then the input data for the task is retrieved from the system memory **230**, according to an embodiment.

(120) Rule 5: if SPL=1, then TQSP is set to 0 for the current task queue **1004**, according to an embodiment. In some embodiments, since Rule 5 sets TQSP to zero, for every task with TSR=1, a task with SPL=1 follows later in the task list **904**.

(121) An example task switch is described with reference to FIG. **15**. FIG. **15** is an example neural network **900** associated with a task queue **1004** that includes tasks T0-T3, according to one embodiment. In this example, the neural task manager **310** directs the neural processor circuit **218** to execute the tasks in order, starting from task T0 and ending with task T3. Additionally, each task can read input data from the previous task. For example, output data **328** for task T0 can be input data for task T1. This can be done by saving the output data **328** of each task to the data buffer **318**. As such, if T0 is the first task in the task list **904**, only T0 will receive input data from the system memory **230** (in absence of a task switch request).

(122) In this example, if a task switch occurs between task T1 and T2. The task switch parameter values can include: T1: TSE=1 DPC=1 TSR=1 T2: SPC=1 SPL=1.

(123) The other task switching parameters may be set to zero. Additional task switch parameter combinations may also allow a task switch between task T1 and T2. For example, T0: TSE=1 T1: DPC=1 TSR=1 T2: SPC=1 SPL=1.

(124) If a task switch request is received while the neural task manager **310** is executing task T1, then the neural task manager **310** can set GTSP to 1, since TSE=1 (see Rule 1). Next, since DPC=1 and the GTSP=1, the neural task manager **310** can direct the output data **328** of T1 to the system memory **230** (see Rule 2). For example, the buffer DMA **320** provides output data **328** from the data buffer **318** to the system memory **230**. Next, since TSR=1 and GTSP=1, the neural task manager **310** can set TQSP to 1 for the interrupted task queue **1004** associated with tasks T0-T3, can set GTSP to 0, and can conduct the task switch to a new task queue **1004** (see Rule 3). In doing this, the configuration data **1014** in the configuration queue **1010** may be replaced with configuration data **1014** associated with a task of a new task queue **1004**, instead of configuration data **1014** associated with a task of the interrupted task queue **1004**.

(125) Since TQSP=1 for the interrupted task queue **1004**, the neural task manager **310** can return to the interrupted task queue **1004** (e.g., after finishing the new task queue **1004**) so that the neural processor circuit **218** can execute tasks T2 and T3. Since TQSP=1 and SPC=1 for task T2, the neural task manager **310** can read the input data from the system memory **230** (see Rule 4). For example, input data may be saved to the data buffer **318** from the system memory **230**. Additionally data may be uploaded from the system memory **230**, such as kernel data and configuration data **1014**. Notice, the input data can be the output data **328** that was previously saved to the system memory **230** when performing Rule 2. Since SPL=1 for task T2, the neural task manager **310** sets TQSP to 0 (see Rule 5). This can complete the task switching process.

(126) Another example of task switching can be described with reference to FIG. **16**. FIG. **16** is an example neural network **900** associated with a task queue **1004** that includes tasks T0-T13,

according to one embodiment. In the example of FIG. 16, the neural task manager **310** directs the neural processor circuit **218** to execute the tasks in order, starting from task T0 and ending on task T13. Tasks T0, T2, T4, and T6, each receive their input data from the system memory **230** (note that the input data from the system memory **230** may be uploaded to the data buffer **318** prior to execution). In this example, a task switch occurs between task T6 and T7. Thus, the output data **328** of tasks T1, T3, T5, and T6 are saved to the system memory **230** and the input data for tasks T7, T9, T11, and T13 is retrieved from the system memory **230**. In some embodiments, input data for tasks T7, T9, T11, and T13 includes the output data **328** of tasks T1, T3, T5, and T6. Tasks T1, T3, T5, T8, T10, and T12 each read input data from a previous task (e.g., output data **328** of task T0 is input data for task T1). This can be done by saving the output data **328** of each task to the data buffer **318** or to system memory **230**.

(127) In this example, a task switch occurs between task T6 and T7. The task switching parameters can include: T1: TSE=1 DPC=1 T3: DPC=1 T5: DPC=1 T6: DPC=1 TSR=1 T7: SPC=1 T9: SPC=1 T11: SPC=1 T13: SPC=1 SPL=1

All other task switching parameters may be set to zero. Additional task switch parameter value combinations may also allow a task switch between task T1 and T2.

(128) Since TSE=1 for T1, if a task switch request is received while the neural task manager **310** is executing task T1, then the neural task manager **310** can set GTSP to 1 (see Rule 1). Next, since DPC=1 for T1, T3, T5, and T6 and GTSP=1, the neural task manager **310** can direct the output data **328** for tasks T1, T3, T5, and T6 to the system memory **230** (see Rule 2). Notice, since DPC=0 for tasks T0, T2, and T4, the output data **328** can be saved to the data buffer **318**. Next, since TSR=1 for T6 and GTSP=1, the neural task manager **310** can set TQSP to 1 for the interrupted task queue **1004** associated with tasks T0-T13, can set GTSP to 0, and can conduct the task switch to a new task queue **1004** (see Rule 3). In these embodiments, note that the task switch does not occur until the neural task manager **310** reads TSR=1 in the task descriptor header **1202** of task T6.

(129) Since TQSP=1 for the interrupted task queue **1004**, the neural task manager **310** can return to the interrupted task queue (e.g., after finishing the new task queue **1004**) so that the neural processor circuit **218** can execute tasks T7-T13. Since TQSP=1 and SPC=1 for tasks T7, T9, T11, and T13, the neural task manager **310** can read the input data from the system memory **230** (see Rule 4). Note, since SPC=0 for tasks T8, T10, and T12, the neural task manager **310** can read the input data from the data buffer **318**. Since SPL=1 for T13, the neural task manager **310** can set TQSP to 0 to complete the task switching process (see Rule 5).

(130) In these embodiments, if no task switch request is received while the neural task manager **310** is executing task T1, then GTSP may not be set to 1 (see Rule 1). As a result, the conditions of Rules 1-5 may not be met, and therefore, no task switch can occur. Additionally, FIG. 16 assumes tasks T7, T9, T11, and T13, use different channels of the inception layer (different parts of the output of T1, T3, T5, and T6). Tasks T7, T9, T11, and T13 can each separately read part of or all the output of T1, T3, T5, and T6. However, if some or all of the tasks have shared input from the previous tasks this input can be read from system memory **230** once and copied to the data buffer **318** for the use on the following tasks. For example, if T7 reads the entire output of T1, T3, T5, and T6, task T7 can have SPC=1 and SPL=1, while the task switching parameters for tasks T9, T11, and T13 can be zero.

(131) The neural task manager **310** may decide which new task queue **1004** to switch to at any time. In some embodiments, the task switch request specifies the new task queue **1004**. In some embodiments, the decision is not made until switching from the interrupted task queue **1004** is complete (e.g., after Rules 1-3 occur). For example, the neural task manager **310**, can assess the priority parameters **1112** of the task queues **1004** after switching from the interrupted task queue **1004**. If two or more task queues **1004** have a same priority parameter **1112**, the neural task manager **318** can select a task queue **1004** by considering additional parameters, such as task queue **1004** indexes.

(132) In some embodiments, a last task in a task list **904** of a current task queue **1004** has TSE=1 and TSR=1. This can allow the neural task manager **310** to switch to a new task queue **1004** after the current task queue **1004** is complete, even without receiving a task switch request. In some embodiments, a first task in a task list **904** of a task queue **1004** may have SPL=1. This can clear any task switching process that may have been set earlier for the task queue **1004**. Furthermore, in embodiments where the first task receives its input data from the external system memory **230**, setting SPL=1 for a first task prevents the first task from unintentionally receiving input data from the data buffer **318**.

(133) FIG. **17** is a flowchart illustrating a method of task switching in the neural processor circuit **218**, according to one embodiment. The method may include different and/or additional steps, or the steps may be in different orders.

(134) The CPU **208** stores **1702** a reference to a first task list **904** of tasks in a first task queue **1004** included in the neural task manager **310**. The first task list **904** of tasks instantiates a first neural network **900** when executed by a neural processor circuit **218**. The first task list **904** may include task descriptors **1012** of tasks stored in the system memory **230**, and the reference may include pointers to memory locations storing the task descriptors **1012**.

(135) The CPU **208** stores **1704** a reference to a second task list **904** of tasks in a second task queue **1004** included in the neural task manager **310**. The second task list **904** of tasks instantiates a second neural network **900** when executed by a neural processor circuit **218**. The second task list **904** may include task descriptors **1012** of tasks stored in the system memory **230**, and the reference may include pointers to memory locations storing the task descriptors **1012**. The first and second task lists **904** may instantiate different neural networks.

(136) The task queues **1004** are a part of the neural task manager **310** and the neural task manager **310** can include any number of task queues **1004** that each stores a reference to a task list **904** of task descriptors **1012** describing tasks. In some embodiments, each of the tasks, when executed, instantiates a single network layer of a neural network **900**, multiple network layers of a neural network **900**, or a portion of a network layer of a neural network **900**. In some embodiments, each task list **904** of task descriptors **1012** includes configuration data **1014** that is stored in a location of the system memory **230**.

(137) The neural task manager **310** (e.g., through the task arbiter **1002**) determines **1706** the priority parameters **1112** of the task queues **1004**. If a task queue **1004** has a higher priority parameter **1112** than the current task queue **1004** being executed, then the neural task manager **310** may initiate a task switch process. For example, if tasks of the first task queue **1004** are being executed by the neural engine **314** and the second task queue **1004** has a higher priority parameter **1112** than the priority parameter **1112** of the first task queue **1004**, then the task manager begins a task switch process. In some embodiments, the neural task manager **310** monitors the the priority parameters **1112** of the task queues **1004** over time to initiate the task switch

(138) During execution of a task of the first task list **904** by the neural engine **314**, the neural task manager **310**, by the task arbiter **1002**, retrieves **1708** configuration data **1014**, for a task in the second task list **904**, from the system memory **230** and sends the configuration data **1014** to the neural processor **218**. For example, the neural processor circuit **218** executes the task of the first task list **904** as discussed above in connection with FIGS. **8** and **14**. While the task of the first task list **904** is executing, the configuration data **1014** programs the neural processor circuit **218** to instantiate a portion of the second neural network **900** by executing the task of the second task list **904**. In some embodiments, the configuration data **1014** for the task in the second task list **904** is retrieved from the system memory **230** using the reference to the second task list **904** stored in the second task queue **1004**.

(139) In some embodiments, the configuration data **1014** for the task in the second task list **904** is stored in a configuration queue **1010** coupled to the neural engine **314** and configured to provide the configuration data **1014** to the neural engine **314**. Furthermore, when the configuration data

**1014** is stored in the configuration queue **1010**, a prefetch process may be initiated where the kernel DMA **324** retrieves kernel data of the task from the system memory **230** to store in the kernel extract **432** and the buffer DMA **320** retrieves input data of the task from the system memory **230** to store in the data buffer **318**.

(140) In some embodiments, the task arbiter **1002** retrieves the configuration data **1014** via the task manager DMA **1006**. Additionally, the configuration data **1014** for the task in the second task list **904** is stored in a fetch queue **1008**. The fetch queue **1008** provides the configuration data **1014** to the configuration queue **1010** after configuration data **1014** previously in the configuration queue **1010** is removed (e.g., after execution of the task defined by the configuration data **1014**). In some embodiments, the configuration data **1014** for the task in the second task list **904** replaces a previously inserted instance of configuration data **1014** in the fetch queue **1008** for the task in the first task list **904**.

(141) The output data **328** of the task in the first task list **904** is stored **1710** the system memory **230**. If a task switch process is not in progress, then the output data **328** may be stored in the data buffer **318** for use as input data in a subsequent processing cycle, and without transfer to the system memory **230**. In some embodiments, after the output data **328** is stored in the system memory **230**, the neural processor circuit **218** begins executing the task of the second task list **904**. One or more tasks in the second task list **904** may be executed by the neural processor circuit **218** to (e.g., fully) instantiate the second neural network **900**.

(142) If the second task list has multiple tasks, the neural engine **314** receives the output data **328** from the executed task in the second task list **904** from the data buffer **318** as input data for a second task of the second task list **904**. From this, the neural engine **314** generates output data **328** from the executed second task of the second task list **904** and provides the output data **328** to the data buffer **318** for another task of the second task list **904**, and so forth until each task in the second task list **904** has been executed. In some embodiments, execution of the tasks of the second task list may also be interrupted with a task switch.

(143) After execution of the tasks in the second task list **904**, the output data **328** of one or more tasks in the second task list **904** are stored **1712** in the system memory **230**. In some embodiments, the priority parameters **1112** of the task queues **1004** continue to be monitored **1706** to determine the next task queue **1004** to execute.

(144) When returning to an unexecuted task of the first task list **904** after execution of the second task list **904**, configuration data **1014** for the unexecuted task is retrieved **1714** from the system memory **230**. The configuration data **1014** may include the output data of the executed task of the first task list **904** stored in the system memory **230** in connection with the task switch. The configuration data **1014** may be used as input data for the next task of the first task list to be executed. The configuration data **1014** may be sent to the neural processor **218** to program the neural processor circuit **218** to instantiate a portion of the first neural network **900** by executing the task of the first task list **904**. In some embodiments, the remaining unexecuted tasks in the first task list **904** are executed by the neural processor circuit **218** to (e.g., fully) instantiate the first neural network **900**.

(145) While particular embodiments and applications have been illustrated and described, it is to be understood that the invention is not limited to the precise construction and components disclosed herein and that various modifications, changes and variations which will be apparent to those skilled in the art may be made in the arrangement, operation and details of the method and apparatus disclosed herein without departing from the spirit and scope of the present disclosure.

## Claims

1. A neural processor circuit, comprising: a neural engine circuit configured to perform neural operations; a neural task manager circuit, comprising: a fetch queue configured to store a first task

descriptor of a task of a first task list in a first task queue, wherein the first task descriptor includes a task descriptor header comprising a task switch parameter defining whether a task switch is to be initiated by the neural task manager circuit; and a task arbiter circuit coupled to the fetch queue and configured to: retrieve the first task descriptor and place the first task descriptor into the fetch queue; and in response to determining that a second task queue has a higher priority than the first task queue, switch to a task of a second task list in the second task queue from the task in the first task queue by replacing the first task descriptor stored in the fetch queue with a second task descriptor of the task in the second task queue; and a configuration queue coupled to the fetch queue and the task arbiter circuit and configured to store configuration data for one or more committed tasks including a committed task committed for execution, wherein the task arbiter circuit is further configured to remove, after the committed task has finished execution, a task descriptor or configuration data of the committed task from the configuration queue.

2. The neural processor circuit of claim 1, further comprising: a task manager direct memory access (DMA) coupled to the fetch queue and the task arbiter circuit, wherein the task arbiter circuit is further configured to: control the task manager DMA to retrieve the first task descriptor or the second task descriptor; and store the first task descriptor or the second task descriptor into the fetch queue.

3. The neural processor circuit of claim 1, wherein the first task queue is coupled to the task arbiter circuit and configured to store a reference to the first task list for instantiating a first neural network by the neural engine circuit configured to perform the neural operations for the task in the first task queue.

4. The neural processor circuit of claim 3, further comprising: a kernel direct memory access (DMA) configured to fetch kernel data from a system memory external to the neural task manager circuit and send the kernel data to the neural engine circuit to perform the neural operations; and a buffer DMA configured to receive input data from the system memory and send the input data to the neural engine circuit to perform the neural operations.

5. The neural processor circuit of claim 4, further comprising: a data buffer configured to store output data of the task in the first task queue used as input data for a subsequent task in the first task queue in response to no task switching for the task in the first task queue, wherein the data buffer is different from the system memory external to the neural task manager circuit.

6. The neural processor circuit of claim 5, wherein the data buffer is coupled to the buffer DMA and configured to store the input data from the system memory before sending the input data to the neural engine circuit to perform the neural operations.

7. The neural processor circuit of claim 1, wherein the first task queue includes a reference to a set of task descriptors including the first task descriptor stored in a system memory external to the neural task manager circuit, a first priority parameter to indicate a first priority for the first task queue, and a task queue switch parameter indicating whether the first task queue was interrupted due to the task switching to the task of the second task queue.

8. The neural processor circuit of claim 1, wherein the first task descriptor includes configuration data of the task in the first task queue, and wherein the task arbiter circuit is configured to extract the configuration data that programs the neural task manager circuit and the neural processor circuit to perform the neural operations.

9. The neural processor circuit of claim 1, wherein the first task descriptor includes the task switch parameter indicating whether the neural task manager circuit is to initiate a task switch after execution of the task in the first task queue.

10. The neural processor circuit of claim 1, wherein the first task descriptor includes an input surface parameter indicating whether input data for the task in the first task queue is to be retrieved from a system memory external to the neural processor circuit or a data buffer included in the neural processor circuit and configured to store output data of the task in the first task queue that is used as input data for a subsequent task in the first task queue.



11. The neural processor circuit of claim 1, wherein the first task descriptor includes an output surface parameter indicating whether output data of the task in the first task queue is to be stored in a system memory external to the neural processor circuit or a data buffer included in the neural processor circuit to facilitate programming of the neural processor circuit.
12. The neural processor circuit of claim 1, wherein the first task descriptor includes a task switch ready (TSR) parameter that defines whether the neural task manager circuit can perform task switching after execution of the task in the first task queue.
13. The neural processor circuit of claim 1, wherein the first task descriptor includes a source pointer last (SPL) parameter that indicates, after returning to an interrupted task queue, the task in the first task queue is a last task with input data stored in a system memory external to the neural processor circuit.
14. A method of task switching in a neural processor circuit, comprising: retrieving a first task descriptor of a task of a first task list in a first task queue including a memory storing a first priority parameter defining a priority of the first task queue; wherein the first task descriptor includes a task descriptor header comprising a task switch parameter defining whether a task switch is to be initiated by a neural task manager circuit; storing the first task descriptor in a fetch queue of the neural task manager circuit, wherein the first task descriptor is placed in a configuration queue for execution of the task in the first task queue; determining, by a task arbiter circuit of the neural task manager circuit, that a second task queue having a second priority parameter that is of a higher priority than the first priority parameter of the first task queue; performing, by the task arbiter circuit, task switching to a task of a second task list in the second task queue from the task in the first task queue by replacing the first task descriptor stored in the fetch queue by a second descriptor of the task in the second task queue; storing, by the configuration queue, configuration data for one or more committed tasks including a committed task committed for execution; and removing, after the committed task has finished execution, a task descriptor or configuration data of the committed task from the configuration queue.
15. The method of claim 14, wherein the first task queue includes a reference to a set of task descriptors including the first task descriptor stored in a system memory external to the neural task manager circuit, the first priority parameter for the first task queue, and a task queue switch parameter indicating whether the first task queue was interrupted due to the task switching to the task of the second task queue.
16. The method of claim 14, wherein the first task descriptor includes configuration data of the task in the first task queue, and wherein the method further comprises extracting, by the task arbiter circuit, the configuration data that programs the neural task manager circuit and the neural processor circuit to perform neural operations.
17. The method of claim 14, wherein the first task descriptor includes the task switch parameter indicating whether the neural task manager circuit is to initiate a task switch after the execution of the task in the first task queue.
18. A system, comprising: a neural engine circuit configured to perform neural operations; a neural task manager circuit, comprising: a fetch queue configured to store a first task descriptor of a task of a first task list in a first task queue including a memory storing a priority parameter defining a priority of the first task queue, wherein the first task descriptor includes a task descriptor header comprising a task switch parameter defining whether a task switch is to be initiated by the neural task manager circuit, wherein the first task descriptor is placed into a configuration queue for execution of the task in the first task queue, and wherein the configuration queue is coupled to the fetch queue and configured to store configuration data for the task in the first task queue once the task in the first task queue has been committed for execution; a task arbiter circuit coupled to the fetch queue and the configuration queue and configured to: retrieve the first task descriptor and place the first task descriptor into the fetch queue; and in response to determining that a second task queue has a higher priority than the first task queue, switch to a task of a second task list in the

second task queue from the task in the first task queue by replacing the first task descriptor stored in the fetch queue with a second task descriptor of a second task; and the configuration queue configured to store configuration data for one or more committed tasks including a committed task committed for execution, wherein the task arbiter circuit is further configured to remove, after the committed task has finished execution, a task descriptor or configuration data of the committed task from the configuration queue.

19. The system of claim 18, further comprising: a task manager direct memory access (DMA) coupled to the fetch queue and the task arbiter circuit, wherein the task arbiter circuit is further configured to: control the task manager DMA to retrieve the first task descriptor of the task in the first task queue or the second task descriptor of the task in the second task queue; and store the first task descriptor or the second task descriptor into the fetch queue.

---