

(12) **United States Patent**
Guim Bernat

(10) **Patent No.:** **US 12,393,456 B2**
(45) **Date of Patent:** **Aug. 19, 2025**

(54) **RESOURCE SELECTION BASED IN PART ON WORKLOAD**

(71) Applicant: **Intel Corporation**, Santa Clara, CA (US)

(72) Inventor: **Francesc Guim Bernat**, Barcelona (ES)

(73) Assignee: **Intel Corporation**, Santa Clara, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 794 days.

(21) Appl. No.: **17/324,525**

(22) Filed: **May 19, 2021**

(65) **Prior Publication Data**
US 2021/0271517 A1 Sep. 2, 2021

(51) **Int. Cl.**
G06F 9/50 (2006.01)
G06N 3/02 (2006.01)
G06N 3/08 (2023.01)

(52) **U.S. Cl.**
CPC **G06F 9/5005** (2013.01); **G06F 9/505** (2013.01); **G06N 3/02** (2013.01); **G06N 3/08** (2013.01)

(58) **Field of Classification Search**
CPC G06F 9/505; G06F 9/5005; G06N 3/08; G06N 3/02
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

9,912,774 B2* 3/2018 Daly H04L 69/22
2016/0379111 A1* 12/2016 Bittner, Jr. G06F 9/3885
706/25

2018/0173291 A1* 6/2018 Levit-Gurevich G06F 1/3206
2019/0042234 A1* 2/2019 Bernat G06F 12/06
2019/0042870 A1* 2/2019 Chen G06F 18/2163
2019/0188570 A1* 6/2019 Aldea Lopez G06N 3/084
2020/0174840 A1* 6/2020 Zhao G06F 9/5011

OTHER PUBLICATIONS

Ullah, Qazi Zia, et al., “Adaptive Resource Utilization Prediction System for Infrastructure as a Service Cloud”, Hindawi, Computational Intelligence and Neuroscience, vol. 2017, Article ID 4873459, Jul. 25, 2017, 12 pages.

* cited by examiner

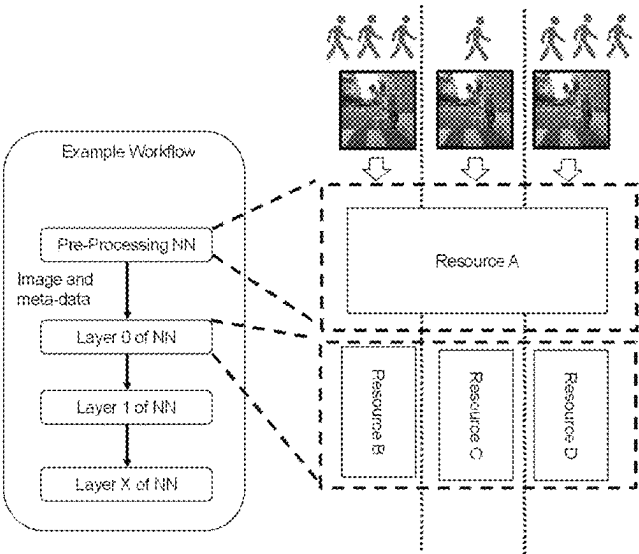
Primary Examiner — Jacob D Dascomb

(74) *Attorney, Agent, or Firm* — Compass IP Law PC

(57) **ABSTRACT**

Examples described herein relate to a system comprising at least one processor and circuitry to: determine multiple configurations of hardware resources to perform a workload associated with a workload request in a subsequent stage based on a pre-processing operation associated with the workload request and at least one service level agreement (SLA) parameter associated with the workload request. In some examples, an executable binary is associated with the workload request and execution of the executable binary performs the pre-processing operation. In some examples, the circuitry is to store the multiple configurations of hardware resources to perform a workload associated with the workload request in a subsequent stage, wherein the multiple configurations of hardware resources are available for access by one or more accelerator devices to perform the workload.

21 Claims, 17 Drawing Sheets






FIG. 1




FIG. 2




FIG. 3




FIG. 4




FIG. 5




FIG. 6




FIG. 7




FIG. 8




FIG. 9




FIG. 10

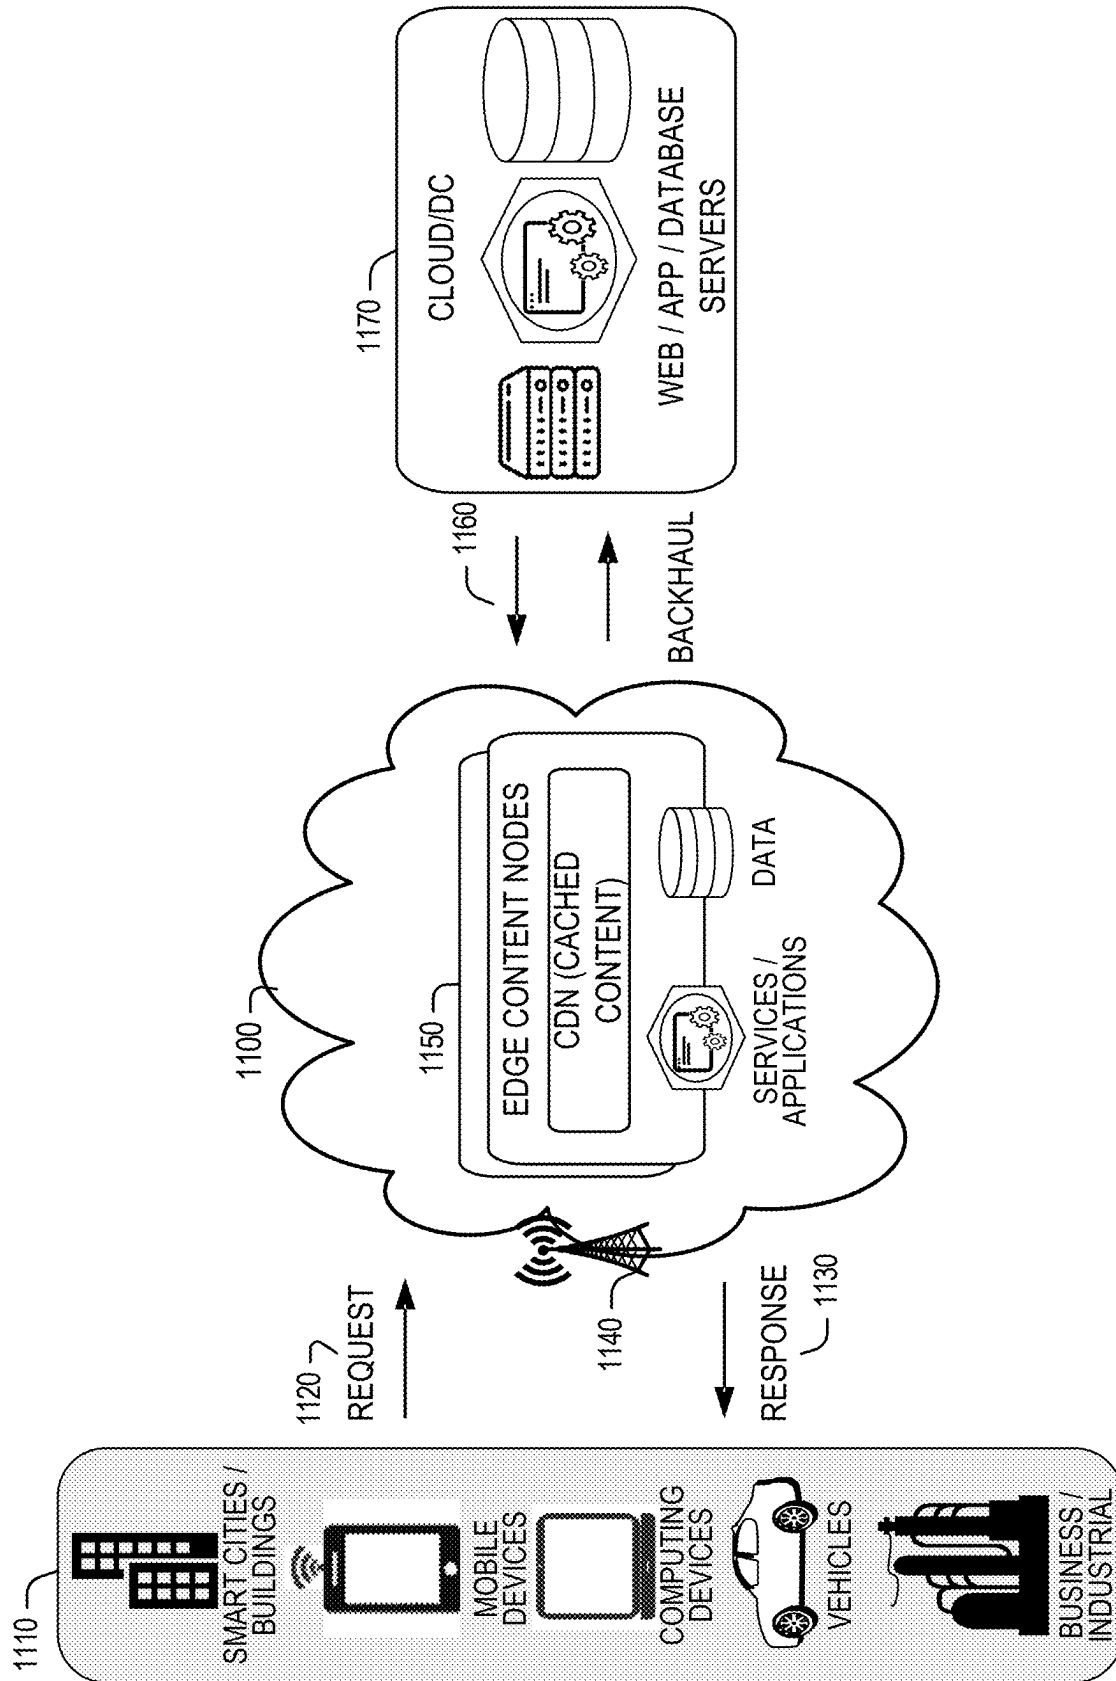


FIG. 11

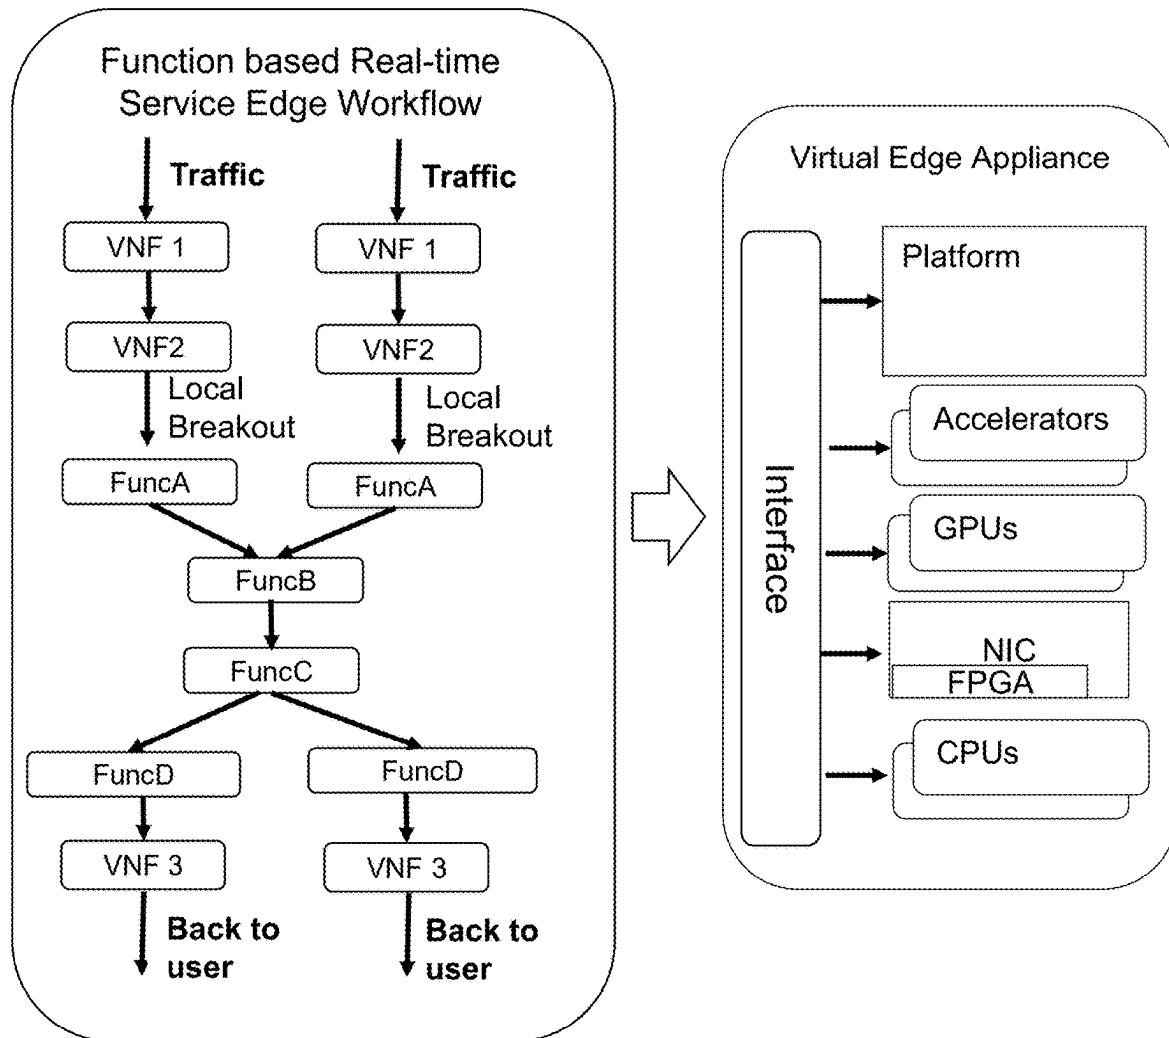


FIG. 12

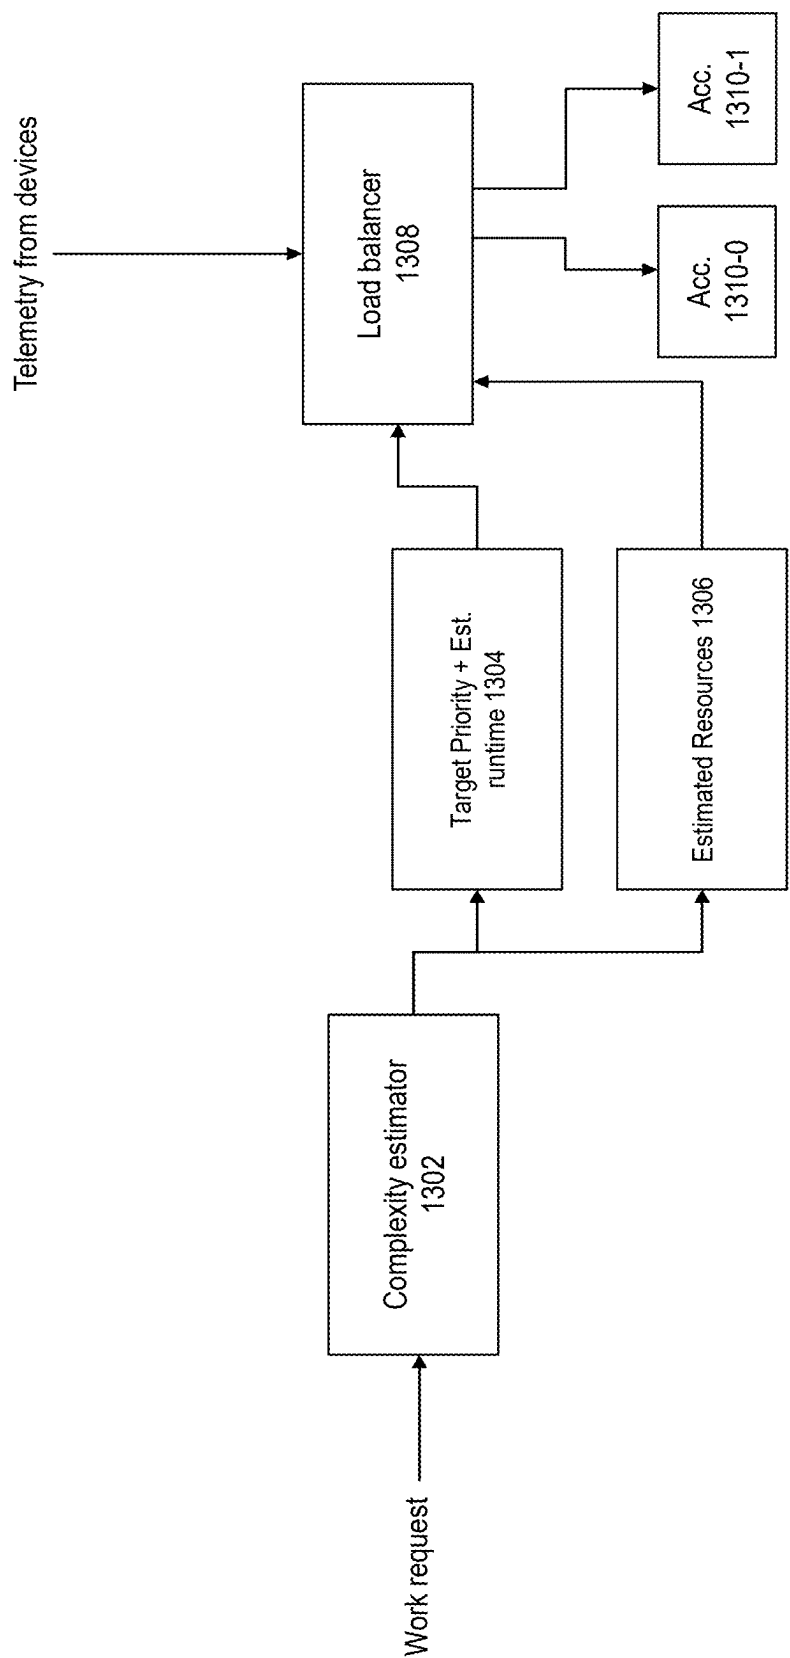


FIG. 13

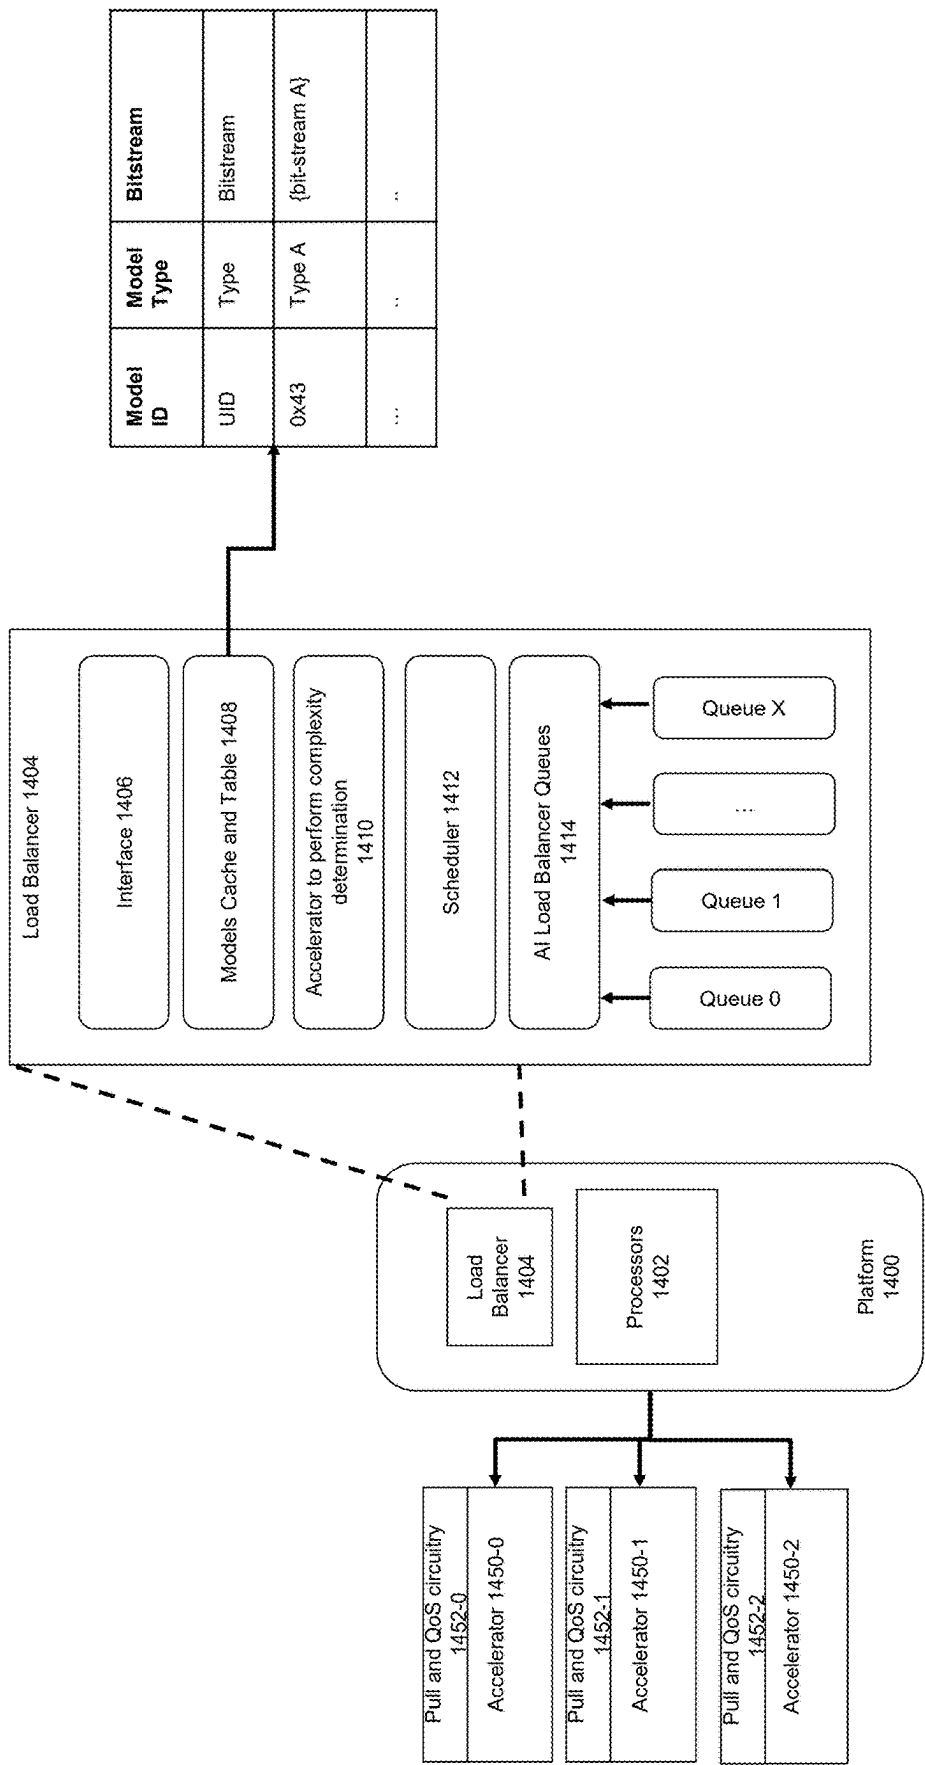


FIG. 14

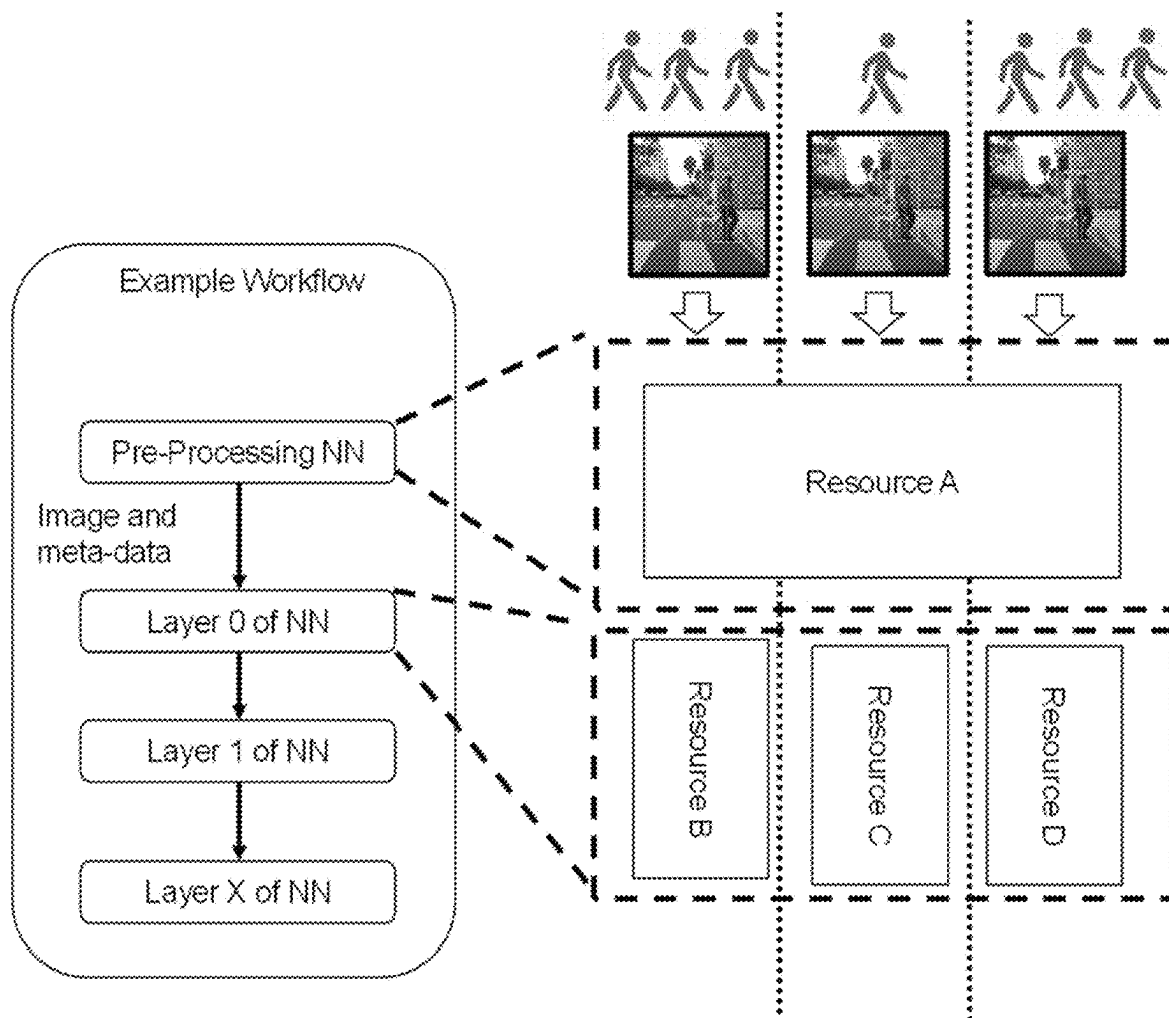


FIG. 15

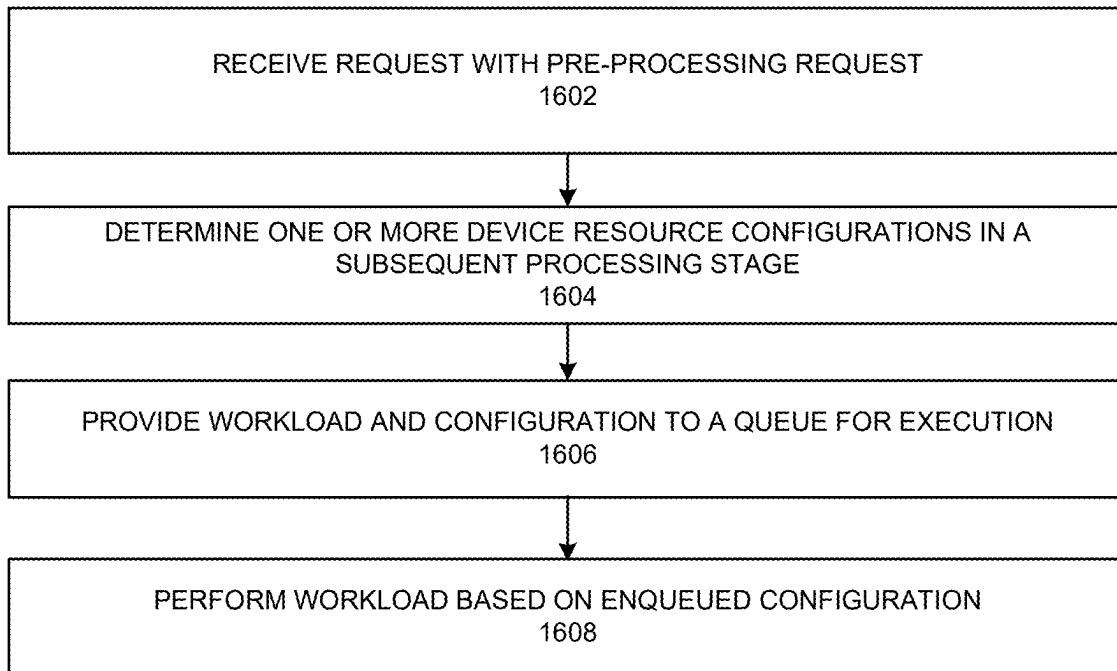


FIG. 16

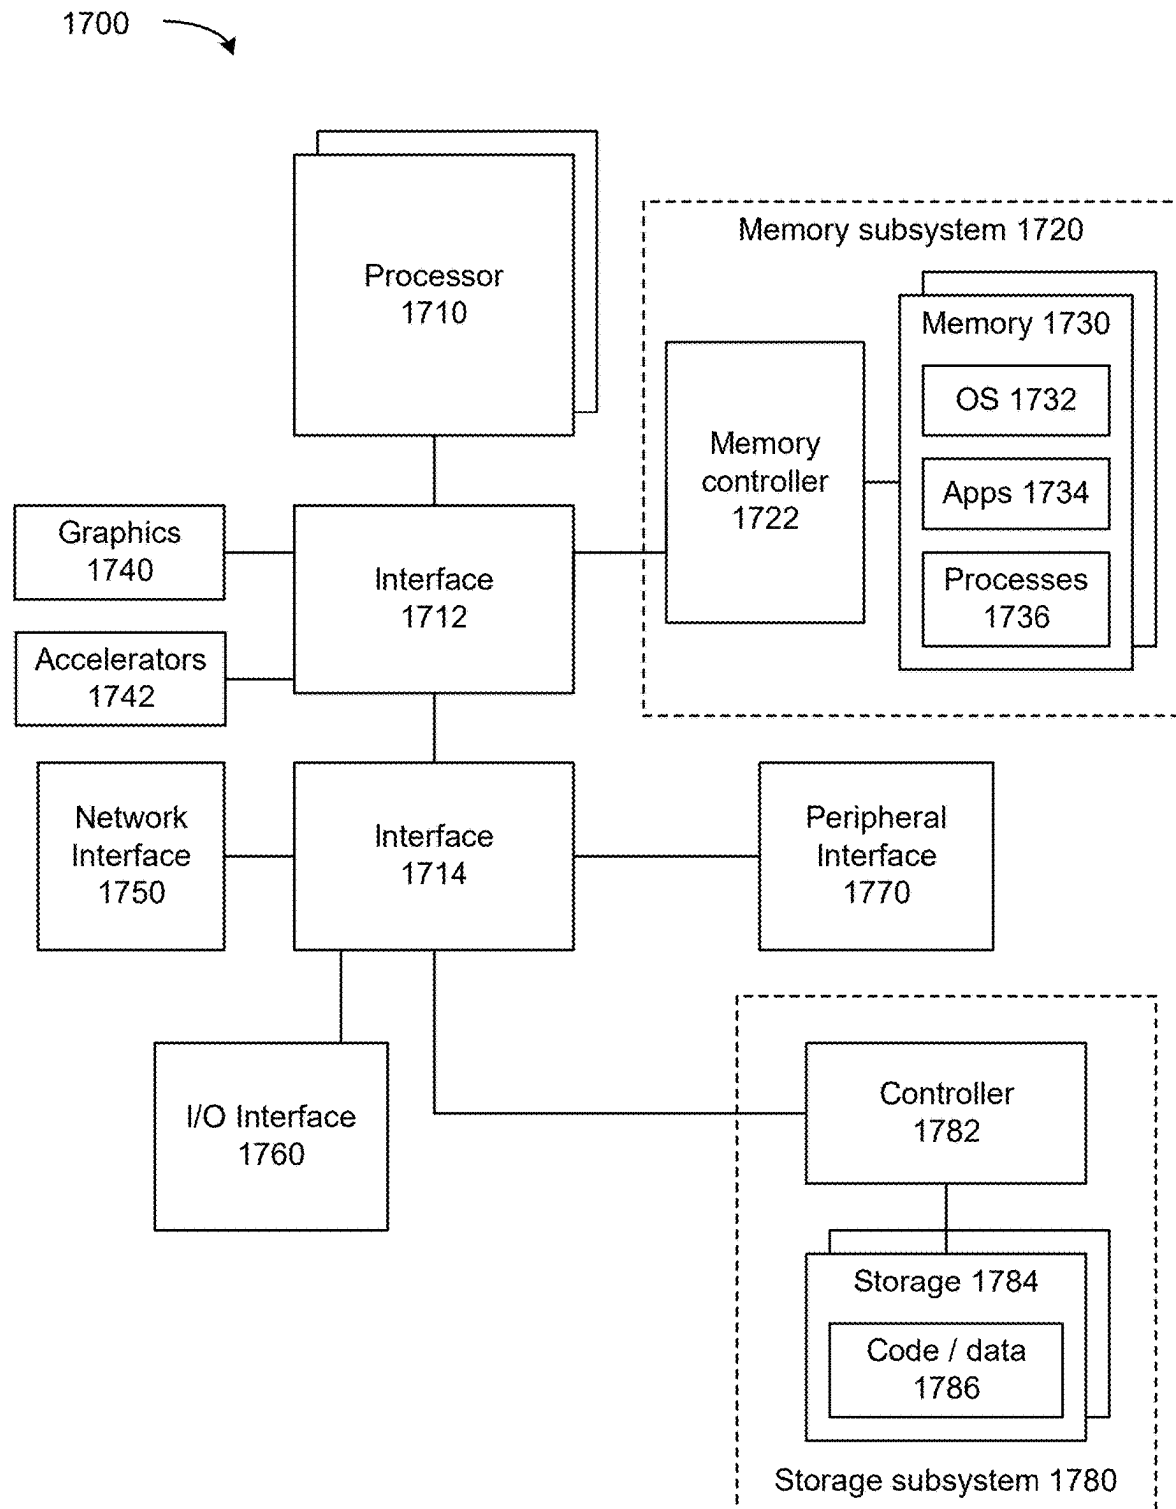


FIG. 17

RESOURCE SELECTION BASED IN PART ON WORKLOAD

BACKGROUND

Cloud computing offers flexibility to select hardware, firmware, and/or software resources. Cloud service providers (CSPs) can offer a platform of resources for use by applications, services, containers, and so forth. Service level agreements (SLAs) can specify resource and performance requirements to perform a workload. CSPs utilize resource management technologies to achieve SLA requirements while balancing resource availability for other uses or workloads.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 is a simplified diagram of at least one data center for executing workloads with disaggregated resources.

FIG. 2 is a simplified diagram of at least one system that may be included in a data center.

FIG. 3 is a simplified block diagram of at least one top side of a node.

FIG. 4 is a simplified block diagram of at least one bottom side of a node.

FIG. 5 is a simplified block diagram of at least one compute node.

FIG. 6 is a simplified block diagram of at least one accelerator node usable in a data center.

FIG. 7 is a simplified block diagram of at least one storage node usable in a data center.

FIG. 8 is a simplified block diagram of at least one memory node usable in a data center.

FIG. 9 depicts a system for executing one or more workloads.

FIG. 10 depicts an example system.

FIG. 11 shows an example system.

FIG. 12 depicts an example edge cloud base station system.

FIG. 13 depicts an example system.

FIG. 14 depicts an example system.

FIG. 15 depicts an example operation of a system.

FIG. 16 depicts an example process.

FIG. 17 depicts an example system.

DETAILED DESCRIPTION

FIG. 1 depicts a data center in which disaggregated resources may cooperatively execute one or more workloads (e.g., applications on behalf of customers) that includes multiple systems **110**, **70**, **130**, **80**, a system being or including one or more rows of racks, racks, or trays. Of course, although data center **100** is shown with multiple systems, the data center **100** may be embodied as a single system. As described in more detail herein, rack houses multiple nodes, some of which may be equipped with one or more type of resources (e.g., memory devices, data storage devices, accelerator devices, general purpose processors, GPUs, xPUs, CPUs, field programmable gate arrays (FPGAs), or application-specific integrated circuits (ASICs)). Resources can be logically coupled or aggregated to form a composed node or composite node, which can act as, for example, a server to perform a job, workload or microservices. One or more examples described herein can be used in a data center or micro data center.

Various examples described herein can perform an application composed of microservices, where a microservice

runs in its own process and communicates using protocols (e.g., application program interface (API), a Hypertext Transfer Protocol (HTTP) resource API, message service, remote procedure calls (RPC), or Google RPC (gRPC)).

Microservices can be independently deployed using centralized management of these services. The management system may be written in different programming languages and use different data storage technologies. A microservice can be characterized by one or more of: use of fine-grained interfaces (to independently deployable services), polyglot programming (e.g., code written in multiple languages to capture additional functionality and efficiency not available in a single language), or lightweight container or virtual machine deployment, and decentralized continuous microservice delivery.

The nodes in system **110**, **70**, **130**, **80** are connected to multiple system switches (e.g., switches that route data communications to and from nodes within the system). Switches can be positioned top of rack (TOR), end of row (EOR), middle of rack (MOR), or a position in a rack or row. The system switches, in turn, connect with spine switches **90** that switch communications among systems (e.g., the systems **110**, **70**, **130**, **80**) in the data center **100**. Nodes may be connected with a fabric using standards described herein or proprietary standards. Nodes may be connected with other fabrics, such as InfiniBand or Ethernet. As described in more detail herein, resources within nodes in the data center **100** may be allocated to a group (referred to herein as a “managed node”) containing resources from one or more nodes to be collectively utilized in the execution of a workload. The workload can execute as if the resources belonging to the managed node were located on the same node. The resources in a managed node may belong to nodes belonging to different racks, and even to different systems **110**, **70**, **130**, **80**. As such, some resources of a single node may be allocated to one managed node while other resources of the same node are allocated to a different managed node (e.g., one processor assigned to one managed node and another processor of the same node assigned to a different managed node).

The disaggregation of resources to nodes comprised predominantly of a single type of resource (e.g., compute nodes comprising primarily compute resources, memory nodes containing primarily memory resources), and the selective allocation and deallocation of the disaggregated resources to form a managed node assigned to execute a workload improves the operation and resource usage of the data center **100** relative to typical data centers comprised of hyperconverged servers containing compute, memory, storage and perhaps additional resources. For example, because nodes predominantly contain resources of a particular type, resources of a given type can be upgraded independently of other resources. Additionally, because different resource types (processors, memory, storage, accelerators, etc.) typically have different refresh rates, greater resource utilization and reduced total cost of ownership may be achieved. For example, a data center operator can upgrade the processors throughout their facility by only swapping out the compute nodes. In such a case, accelerator and storage resources may not be contemporaneously upgraded and, rather, may be allowed to continue operating until those resources are scheduled for their own refresh. Resource utilization may also increase. For example, if managed nodes are composed based on requirements of the workloads that will be running on them, resources within a node are more likely to be fully utilized. Such utilization may allow for more managed nodes to run in a data center with a given set of resources,

or for a data center expected to run a given set of workloads, to be built using fewer resources.

FIG. 2 depicts a system. A system can include a set of rows **200**, **210**, **220**, **230** of racks **240**. Rack **240** may house multiple nodes (e.g., sixteen nodes) and provide power and data connections to the housed nodes, as described in more detail herein. Racks in row **200**, **210**, **220**, **230** are connected to multiple system switches **250**, **260**. The system switch **250** includes a set of ports **252** to which the nodes of the racks of the system **110** are connected and another set of ports **254** that connect the system **110** to the spine switches **90** to provide connectivity to other systems in the data center **100**. Similarly, the system switch **260** includes a set of ports **262** to which the nodes of the racks of the system **110** are connected and a set of ports **264** that connect the system **110** to the spine switches **90**. As such, the use of the pair of switches **250**, **260** provides an amount of redundancy to the system **110**. For example, if either of the switches **250**, **260** fails, the nodes in the system **110** may still maintain data communication with the remainder of the data center **100** (e.g., nodes of other systems) through the other switch **250**, **260**. Furthermore, switches **90**, **250**, **260** may be embodied as dual-mode optical switches, capable of routing both Ethernet protocol communications carrying Internet Protocol (IP) packets and communications according to a second, high-performance link-layer protocol (e.g., PCI Express or Compute Express Link) via optical signaling media of an optical fabric.

It should be appreciated that systems **70**, **130**, **80** (as well as additional systems of the data center **100**) may be similarly structured as, and have components similar to, the system **110** shown in and described in regard to FIG. 2 (e.g., system may have rows of racks housing multiple nodes as described above). Additionally, while two system switches **250**, **260** are shown, it should be understood, system **110**, **70**, **130**, **80** may be connected to a different number of system switches, providing even more failover capacity. Of course, systems may be arranged differently than the rows-of-racks configuration shown in FIGS. 1-2. For example, a system may be embodied as multiple sets of racks in which set of racks is arranged radially, e.g., the racks are equidistant from a center switch.

Referring now to FIG. 3, node **400** may be configured to be mounted in a corresponding rack **240** of the data center **100** as discussed above. A node **400** may be optimized or otherwise configured for performing particular tasks, such as compute tasks, acceleration tasks, data storage tasks, etc. For example, the node **400** may be embodied as a compute node **500** as discussed below in regard to FIG. 5, an accelerator node **600** as discussed below in regard to FIG. 6, a storage node **700** as discussed below in regard to FIG. 7, or as a node optimized or otherwise configured to perform other specialized tasks, such as a memory node **800**, discussed below in regard to FIG. 8.

Although two physical resources **320** are shown in FIG. 3, it should be appreciated that the node **400** may include one, two, or more physical resources **320**. The physical resources **320** may be embodied as any type of processor, controller, or other compute circuit capable of performing various tasks such as compute functions and/or controlling the functions of the node **400** depending on, for example, the type or intended functionality of the node **400**. For example, as discussed in more detail below, the physical resources **320** may be embodied as high-performance processors in which the node **400** is embodied as a compute node, as accelerator co-processors or circuits in which the node **400** is embodied as an accelerator node, storage controllers in which the node

400 is embodied as a storage node, or a set of memory devices in which the node **400** is embodied as a memory node.

The node **400** also includes one or more additional physical resources **330** mounted to circuit board substrate **302**. The additional physical resources include a network interface controller (NIC) as discussed in more detail below. Of course, depending on the type and functionality of the node **400**, the physical resources **330** may include additional or other electrical components, circuits, and/or devices.

The physical resources **320** can be communicatively coupled to the physical resources **330** via an input/output (I/O) subsystem **322**. The I/O subsystem **322** may be embodied as circuitry and/or components to facilitate input/output operations with the physical resources **320**, the physical resources **330**, and/or other components of the node **400**. For example, the I/O subsystem **322** may be embodied as, or otherwise include, memory controller hubs, input/output control hubs, integrated sensor hubs, firmware devices, communication links (e.g., point-to-point links, bus links, wires, cables, waveguides, light guides, printed circuit board traces, etc.), and/or other components and subsystems to facilitate the input/output operations. The I/O subsystem **322** is embodied as, or otherwise includes, a double data rate 4 (DDR4) data bus or a DDR5 data bus.

The node **400** may also include a resource-to-resource interconnect **324**. The resource-to-resource interconnect **324** may be embodied as any type of communication interconnect capable of facilitating resource-to-resource communications. The resource-to-resource interconnect **324** is embodied as a high-speed point-to-point interconnect (e.g., faster than the I/O subsystem **322**). For example, the resource-to-resource interconnect **324** may be embodied as a QuickPath Interconnect (QPI), an UltraPath Interconnect (UPI), PCI express (PCIe), or other high-speed point-to-point interconnect dedicated to resource-to-resource communications.

The node **400** also includes a power connector **340** configured to mate with a corresponding power connector of the rack **240** when the node **400** is mounted in the corresponding rack **240**. The node **400** receives power from a power supply of the rack **240** via the power connector **340** to supply power to the various electrical components of the node **400**. In some examples, the node **400** includes local power supply (e.g., an on-board power supply) to provide power to the electrical components of the node **400**. In some examples, the node **400** does not include any local power supply (e.g., an on-board power supply) to provide power to the electrical components of the node **400**. The exclusion of a local or on-board power supply facilitates the reduction in the overall footprint of the circuit board substrate **302**, which may increase the thermal cooling characteristics of the various electrical components mounted on the circuit board substrate **302** as discussed above. Voltage regulators are placed on circuit board substrate **302** directly opposite of the processors **520** (see FIG. 5), and power is routed from the voltage regulators to the processors **520** by vias extending through the circuit board substrate **302**. Such a configuration provides an increased thermal budget, additional current and/or voltage, and better voltage control relative to typical printed circuit boards in which processor power is delivered from a voltage regulator, in part, by printed circuit traces.

The node **400** may also include mounting features **342** configured to mate with a mounting arm, or other structure, of a robot to facilitate the placement of the node **300** in a rack **240** by the robot. The mounting features **342** may be embodied as any type of physical structures that allow the

5

robot to grasp the node **400** without damaging the circuit board substrate **302** or the electrical components mounted thereto.

Referring now to FIG. 4, in addition to the physical resources **330** mounted on circuit board substrate **302**, the node **400** also includes one or more memory devices **420**. The physical resources **320** can be communicatively coupled to memory devices **420** via the I/O subsystem **322**. For example, the physical resources **320** and the memory devices **420** may be communicatively coupled by one or more vias extending through the circuit board substrate **302**. A physical resource **320** may be communicatively coupled to a different set of one or more memory devices **420**. Alternatively, physical resource **320** may be communicatively coupled to memory device **420**.

The memory devices **420** may be embodied as any type of memory device capable of storing data for the physical resources **320** during operation of the node **400**, such as any type of volatile (e.g., dynamic random access memory (DRAM), etc.) or non-volatile memory. Volatile memory may be a storage medium that requires power to maintain the state of data stored by the medium. Non-limiting examples of volatile memory may include various types of random access memory (RAM), such as dynamic random access memory (DRAM) or static random access memory (SRAM). One particular type of DRAM that may be used in a memory module is synchronous dynamic random access memory (SDRAM). DRAM of a memory component may comply with a standard promulgated by JEDEC, such as JESD79F for DDR SDRAM, JESD79-2F for DDR2 SDRAM, JESD79-3F for DDR3 SDRAM, JESD79-4A for DDR4 SDRAM, JESD209 for Low Power DDR (LPDDR), JESD209-2 for LPDDR2, JESD209-3 for LPDDR3, and JESD209-4 for LPDDR4. Such standards (and similar standards) may be referred to as DDR-based standards and communication interfaces of the storage devices that implement such standards may be referred to as DDR-based interfaces.

Memory device can be block addressable memory device, such as those based on NAND or NOR technologies, for example, multi-threshold level NAND flash memory and NOR flash memory. A block can be any size such as but not limited to 2 KB, 4 KB, 5 KB, and so forth. A memory device may also include next-generation nonvolatile devices, such as Intel Optane® memory or other byte addressable write-in-place nonvolatile memory devices (e.g., memory devices that use chalcogenide glass), multi-threshold level NAND flash memory, NOR flash memory, single or multi-level Phase Change Memory (PCM), a resistive memory, nanowire memory, ferroelectric transistor random access memory (FeTRAM), anti-ferroelectric memory, magnetoresistive random access memory (MRAM) memory that incorporates memristor technology, resistive memory including the metal oxide base, the oxygen vacancy base and the conductive bridge Random Access Memory (CB-RAM), or spin transfer torque (STT)-MRAM, a spintronic magnetic junction memory based device, a magnetic tunneling junction (MTJ) based device, a DW (Domain Wall) and SOT (Spin Orbit Transfer) based device, a thyristor based memory device, or a combination of one or more of the above, or other memory. The memory device may refer to the die itself and/or to a packaged memory product. The memory device may comprise a transistor-less stackable cross point architecture in which memory cells sit at the intersection of word lines and bit lines and are individually addressable and in which bit storage is based on a change in bulk resistance.

6

Referring now to FIG. 5, the node **400** may be embodied as a compute node **500**. The compute node **500** can be configured to perform compute tasks. Of course, as discussed above, the compute node **500** may rely on other nodes, such as acceleration nodes and/or storage nodes, to perform compute tasks. In the illustrative compute node **500**, the physical resources **320** are embodied as processors **520**. Although only two processors **520** are shown in FIG. 5, it should be appreciated that the compute node **500** may include additional processors **520**. Illustratively, the processors **520** are embodied as high-performance processors **520** and may be configured to operate at a relatively high power rating.

The compute node **500** may also include a processor-to-processor interconnect **542**. Processor-to-processor interconnect **542** may be embodied as any type of communication interconnect capable of facilitating processor-to-processor interconnect **542** communications. The processor-to-processor interconnect **542** is embodied as a high-speed point-to-point interconnect (e.g., faster than the I/O subsystem **322**). For example, the processor-to-processor interconnect **542** may be embodied as a QuickPath Interconnect (QPI), an UltraPath Interconnect (UPI), or other high-speed point-to-point interconnect dedicated to processor-to-processor communications (e.g., PCIe or Compute Express Link (CXL)).

The compute node **500** also includes a communication circuit **530**. The illustrative communication circuit **530** includes a network interface controller (NIC) **532**, which may also be referred to as a host fabric interface (HFI). The NIC **532** may be embodied as, or otherwise include, any type of integrated circuit, discrete circuits, controller chips, chipsets, add-in-boards, daughtercards, network interface cards, or other devices that may be used by the compute node **500** to connect with another compute device (e.g., with other nodes **400**). NIC **532** may be embodied as part of a system-on-a-chip (SoC) that includes one or more processors, or included on a multichip package that also contains one or more processors. NIC **532** may include a local processor (not shown) and/or a local memory (not shown) that are both local to the NIC **532**. The local processor of the NIC **532** may be capable of performing one or more of the functions of the processors **520**. Additionally or alternatively, the local memory of the NIC **532** may be integrated into one or more components of the compute node at the board level, socket level, chip level, and/or other levels. In some examples, a network interface includes a network interface controller or a network interface card. In some examples, a network interface can include one or more of a network interface controller (NIC) **532**, a host fabric interface (HFI), a host bus adapter (HBA), network interface connected to a bus or connection (e.g., PCIe, CXL, DDR, and so forth). In some examples, a network interface can be part of a switch or a system-on-chip (SoC).

Some examples of a NIC are part of an Infrastructure Processing Unit (IPU) or data processing unit (DPU) or utilized by an IPU or DPU. An IPU or DPU can include a network interface, memory devices, and one or more programmable or fixed function processors (e.g., CPU or XPU) to perform offload of operations that could have been performed by a host CPU, GPU, GPGPU, or XPU or remote CPU or XPU. In some examples, the IPU or DPU can perform virtual switch operations, manage storage transactions (e.g., compression, cryptography, virtualization), and manage operations performed on other IPUs, DPUs, servers, or devices.

The communication circuit **530** is communicatively coupled to an optical data connector **534**. The optical data connector **534** is configured to mate with a corresponding optical data connector of a rack when the compute node **500** is mounted in the rack. Illustratively, the optical data connector **534** includes a plurality of optical fibers which lead from a mating surface of the optical data connector **534** to an optical transceiver **536**. The optical transceiver **536** is configured to convert incoming optical signals from the rack-side optical data connector to electrical signals and to convert electrical signals to outgoing optical signals to the rack-side optical data connector. Although shown as forming part of the optical data connector **534**, the optical transceiver **536** may form a portion of the communication circuit **530**.

Compute node **500** may also include an expansion connector **540**. Expansion connector **540** is configured to mate with a corresponding connector of an expansion circuit board substrate to provide additional physical resources to the compute node **500**. The additional physical resources may be used, for example, by the processors **520** during operation of the compute node **500**. The expansion circuit board substrate may be substantially similar to the circuit board substrate **302** discussed above and may include various electrical components mounted thereto. The particular electrical components mounted to the expansion circuit board substrate may depend on the intended functionality of the expansion circuit board substrate. For example, the expansion circuit board substrate may provide additional compute resources, memory resources, and/or storage resources. As such, the additional physical resources of the expansion circuit board substrate may include, but is not limited to, processors, memory devices, storage devices, and/or accelerator circuits including, for example, field programmable gate arrays (FPGA), application-specific integrated circuits (ASICs), security co-processors, graphics processing units (GPUs), machine learning circuits, or other specialized processors, controllers, devices, and/or circuits. Note that reference to GPU or CPU herein can in addition or alternatively refer to an XPU or xPU. An xPU can include one or more of: a GPU, ASIC, FPGA, or accelerator device.

Referring now to FIG. 6, the node **400** may be embodied as an accelerator node **600**. The accelerator node **600** is configured to perform specialized compute tasks, such as machine learning, encryption, hashing, or other computational-intensive task. For example, a compute node **500** may offload tasks to the accelerator node **600** during operation. The accelerator node **600** includes various components similar to components of the node **400** and/or compute node **500**, which have been identified in FIG. 6 using the same reference numbers.

In the illustrative accelerator node **600**, the physical resources **320** are embodied as accelerator circuits **620**. Although only two accelerator circuits **620** are shown in FIG. 6, it should be appreciated that the accelerator node **600** may include additional accelerator circuits **620**. The accelerator circuits **620** may be embodied as any type of processor, co-processor, compute circuit, or other device capable of performing compute or processing operations. For example, the accelerator circuits **620** may be embodied as, for example, central processing units, cores, field programmable gate arrays (FPGA), application-specific integrated circuits (ASICs), programmable control logic (PCL), security co-processors, graphics processing units (GPUs), neuromorphic processor units, quantum computers, machine learning circuits, programmable processing pipeline (e.g., programmable by P4, C, Python, Broadcom Network Programming Language (NPL), or x86 compatible executable binaries or

other executable binaries). Processors, FPGAs, other specialized processors, controllers, devices, and/or circuits can be used utilized for packet processing or packet modification. Ternary content-addressable memory (TCAM) can be used for parallel match-action or look-up operations on packet header content.

The accelerator node **600** may also include an accelerator-to-accelerator interconnect **642**. Similar to the resource-to-resource interconnect **324** of the node **300** discussed above, the accelerator-to-accelerator interconnect **642** may be embodied as any type of communication interconnect capable of facilitating accelerator-to-accelerator communications. The accelerator-to-accelerator interconnect **642** can be embodied as a high-speed point-to-point interconnect (e.g., faster than the I/O subsystem **322**). For example, the accelerator-to-accelerator interconnect **642** may be embodied as a QuickPath Interconnect (QPI), an UltraPath Interconnect (UPI), or other high-speed point-to-point interconnect dedicated to processor-to-processor communications. The accelerator circuits **620** may be daisy-chained with a primary accelerator circuit **620** connected to the NIC **532** and memory **420** through the I/O subsystem **322** and a secondary accelerator circuit **620** connected to the NIC **532** and memory **420** through a primary accelerator circuit **620**.

Referring now to FIG. 7, node **400** may be embodied as a storage node **700**. The storage node **700** is configured, to store data in a data storage **750** local to the storage node **700**. For example, during operation, a compute node **500** or an accelerator node **600** may store and retrieve data from the data storage **750** of the storage node **700**. The storage node **700** includes various components similar to components of the node **400** and/or the compute node **500**, which have been identified in FIG. 7 using the same reference numbers.

In the illustrative storage node **700**, the physical resources **320** are embodied as storage controllers **720**. Although only two storage controllers **720** are shown in FIG. 7, it should be appreciated that the storage node **700** may include additional storage controllers **720**. The storage controllers **720** may be embodied as any type of processor, controller, or control circuit capable of controlling the storage and retrieval of data into the data storage **750** based on requests received via the communication circuit **530**. The storage controllers **720** are embodied as relatively low-power processors or controllers.

The storage node **700** may also include a controller-to-controller interconnect **742**. Similar to the resource-to-resource interconnect **324** of the node **400** discussed above, the controller-to-controller interconnect **742** may be embodied as any type of communication interconnect capable of facilitating controller-to-controller communications. The controller-to-controller interconnect **742** is embodied as a high-speed point-to-point interconnect (e.g., faster than the I/O subsystem **322**). For example, the controller-to-controller interconnect **742** may be embodied as a QuickPath Interconnect (QPI), an UltraPath Interconnect (UPI), or other high-speed point-to-point interconnect dedicated to processor-to-processor communications.

Referring now to FIG. 8, the node **400** may be embodied as a memory node **800**. The memory node **800** is configured to provide other nodes **400** (e.g., compute nodes **500**, accelerator nodes **600**, etc.) with access to a pool of memory (e.g., in two or more sets **830**, **832** of memory devices **420**) local to the storage node **700**. For example, during operation, a compute node **500** or an accelerator node **600** may remotely write to and/or read from one or more of the

memory sets **830, 832** of the memory node **800** using a logical address space that maps to physical addresses in the memory sets **830, 832**.

In the illustrative memory node **800**, the physical resources **320** are embodied as memory controllers **820**. Although only two memory controllers **820** are shown in FIG. **8**, it should be appreciated that the memory node **800** may include additional memory controllers **820**. The memory controllers **820** may be embodied as any type of processor, controller, or control circuit capable of controlling the writing and reading of data into the memory sets **830, 832** based on requests received via the communication circuit **530**. Memory controller **820** is connected to a corresponding memory set **830, 832** to write to and read from memory devices **420** within the corresponding memory set **830, 832** and enforce a permissions (e.g., read, write, etc.) associated with node **400** that has sent a request to the memory node **800** to perform a memory access operation (e.g., read or write).

The memory node **800** may also include a controller-to-controller interconnect **842**. Similar to the resource-to-resource interconnect **324** of the node **400** discussed above, the controller-to-controller interconnect **842** may be embodied as any type of communication interconnect capable of facilitating controller-to-controller communications. The controller-to-controller interconnect **842** is embodied as a high-speed point-to-point interconnect (e.g., faster than the I/O subsystem **322**). For example, the controller-to-controller interconnect **842** may be embodied as a QuickPath Interconnect (QPI), an UltraPath Interconnect (UPI), or other high-speed point-to-point interconnect dedicated to processor-to-processor communications. As such, a memory controller **820** may access, through the controller-to-controller interconnect **842**, memory that is within the memory set **832** associated with another memory controller **820**. A scalable memory controller is made of multiple smaller memory controllers, referred to herein as “chiplets”, on a memory node (e.g., the memory node **800**). The chiplets may be interconnected (e.g., using EMIB (Embedded Multi-Die Interconnect Bridge)). The combined chiplet memory controller may scale up to a relatively large number of memory controllers and I/O ports, (e.g., up to 16 memory channels). Memory controllers **820** may implement a memory interleave (e.g., one memory address is mapped to the memory set **830**, the next memory address is mapped to the memory set **832**, and the third address is mapped to the memory set **830**, etc.). The interleaving may be managed within the memory controllers **820**, or from CPU sockets (e.g., of the compute node **500**) across network links to the memory sets **830, 832**, and may improve the latency associated with performing memory access operations as compared to accessing contiguous memory addresses from the same memory device.

Memory node **800** may be connected to one or more other nodes **400** (e.g., in the same rack **240** or an adjacent rack **240**) through a waveguide, using the waveguide connector **880**. Using a waveguide may provide high throughput access to the memory pool (e.g., the memory sets **830, 832**) to another node (e.g., a node **400** in the same rack **240** or an adjacent rack **240** as the memory node **800**) without adding to the load on the optical data connector **534**.

Referring now to FIG. **9**, a system for executing one or more workloads (e.g., applications) may be implemented. System **910** includes an orchestrator server **920**, which may be embodied as a managed node comprising a compute device (e.g., a processor **520** on a compute node **500**) executing management software (e.g., a cloud operating

environment, such as OpenStack) that is communicatively coupled to multiple nodes **400** including a large number of compute nodes **930** (e.g., one or more similar to the compute node **500**), memory nodes **940** (e.g., one or more similar to the memory node **800**), accelerator nodes **950** (e.g., one or more similar to the memory node **600**), and storage nodes **960** (e.g., one or more similar to the storage node **700**). One or more of the nodes **930, 940, 950, 960** may be grouped into a managed node **970**, such as by the orchestrator server **920**, to collectively perform a workload (e.g., an application **932** executed in a virtual machine or in a container).

The managed node **970** may be embodied as an assembly of physical resources **320**, such as processors **520**, memory resources **420**, accelerator circuits **620**, or data storage **750**, from the same or different nodes **400**. Further, the managed node may be established, defined, or “spun up” by the orchestrator server **920** at the time a workload is to be assigned to the managed node or at a time, and may exist regardless of whether a workload is presently assigned to the managed node. The orchestrator server **920** may selectively allocate and/or deallocate physical resources **320** from the nodes **400** and/or add or remove one or more nodes **400** from the managed node **970** as a function of quality of service (QoS) targets (e.g., a target throughput, a target latency, a target number instructions per second, etc.) associated with a service level agreement or class of service (COS or CLOS) for the workload (e.g., the application **932**). In doing so, the orchestrator server **920** may receive telemetry data indicative of performance conditions (e.g., throughput, latency, instructions per second, etc.) in a node **400** of the managed node **970** and compare the telemetry data to the quality of service targets to determine whether the quality of service targets are being satisfied. The orchestrator server **920** may additionally determine whether one or more physical resources may be deallocated from the managed node **970** while still satisfying the QoS targets, thereby freeing up those physical resources for use in another managed node (e.g., to execute a different workload). Alternatively, if the QoS targets are not presently satisfied, the orchestrator server **920** may determine to dynamically allocate additional physical resources to assist in the execution of the workload (e.g., the application **932**) while the workload is executing. Similarly, the orchestrator server **920** may determine to dynamically deallocate physical resources from a managed node if the orchestrator server **920** determines that deallocating the physical resource would result in QoS targets still being met.

Additionally, the orchestrator server **920** may identify trends in the resource utilization of the workload (e.g., the application **932**), such as by identifying phases of execution (e.g., time periods in which different operations, one or more having different resource utilizations characteristics, are performed) of the workload (e.g., the application **932**) and pre-emptively identifying available resources in the data center and allocating them to the managed node **970** (e.g., within a predefined time period of the associated phase beginning). The orchestrator server **920** may model performance based on various latencies and a distribution scheme to place workloads among compute nodes and other resources (e.g., accelerator nodes, memory nodes, storage nodes) in the data center. For example, the orchestrator server **920** may utilize a model that accounts for the performance of resources on the nodes **400** (e.g., FPGA performance, memory access latency, etc.) and the performance (e.g., congestion, latency, bandwidth) of the path through the network to the resource (e.g., FPGA). As such, the orchestrator server **920** may determine which resource(s) should be

used with which workloads based on the total latency associated with one or more potential resource available in the data center **100** (e.g., the latency associated with the performance of the resource itself in addition to the latency associated with the path through the network between the compute node executing the workload and the node **400** on which the resource is located).

The orchestrator server **920** may generate a map of heat generation in the data center **100** using telemetry data (e.g., temperatures, fan speeds, etc.) reported from the nodes **400** and allocate resources to managed nodes as a function of the map of heat generation and predicted heat generation associated with different workloads, to maintain a target temperature and heat distribution in the data center **100**. Additionally or alternatively, the orchestrator server **920** may organize received telemetry data into a hierarchical model that is indicative of a relationship between the managed nodes (e.g., a spatial relationship such as the physical locations of the resources of the managed nodes within the data center **100** and/or a functional relationship, such as groupings of the managed nodes by the customers the managed nodes provide services for, the types of functions typically performed by the managed nodes, managed nodes that typically share or exchange workloads among others, etc.). Based on differences in the physical locations and resources in the managed nodes, a given workload may exhibit different resource utilizations (e.g., cause a different internal temperature, use a different percentage of processor or memory capacity) across the resources of different managed nodes. The orchestrator server **920** may determine the differences based on the telemetry data stored in the hierarchical model and factor the differences into a prediction of future resource utilization of a workload if the workload is reassigned from one managed node to another managed node, to accurately balance resource utilization in the data center **100**. The orchestrator server **920** may identify patterns in resource utilization phases of the workloads and use the patterns to predict future resource utilization of the workloads.

To reduce the computational load on the orchestrator server **920** and the data transfer load on the network, the orchestrator server **920** may send self-test information to the nodes **400** to enable a node **400** to locally (e.g., on the node **400**) determine whether telemetry data generated by the node **400** satisfies one or more conditions (e.g., an available capacity that satisfies a predefined threshold, a temperature that satisfies a predefined threshold, etc.). Node **400** may then report back a simplified result (e.g., yes or no) to the orchestrator server **920**, which the orchestrator server **920** may utilize in determining the allocation of resources to managed nodes.

Examples described herein can be used in a data center or disaggregated composite nodes. The techniques described herein can apply to both disaggregated and traditional server architectures. A traditional server can include a CPU, XPU, one or more memory devices, networking communicatively coupled to one or more circuit boards within a server.

Edge Network

Edge computing, at a general level, refers to the implementation, coordination, and use of computing and resources at locations closer to the “edge” or collection of “edges” of the network. The purpose of this arrangement is to improve total cost of ownership, reduce application and network latency, reduce network backhaul traffic and associated energy consumption, improve service capabilities, and improve compliance with security or data privacy requirements (especially as compared to conventional cloud

computing). Components that can perform edge computing operations (“edge nodes”) can reside in whatever location needed by the system architecture or ad hoc service (e.g., in a high performance compute data center or cloud installation; a designated edge node server, an enterprise server, a roadside server, a telecom central office; or a local or peer at-the-edge device being served consuming edge services).

With the illustrative edge networking systems described below, computing and storage resources are moved closer to the edge of the network (e.g., closer to the clients, endpoint devices, or “things”). By moving the computing and storage resources closer to the device producing or using the data, various latency, compliance, and/or monetary or resource cost constraints may be achievable relative to a standard networked (e.g., cloud computing) system. To do so, in some examples, pools of compute, memory, and/or storage resources may be located in, or otherwise equipped with, local servers, routers, and/or other network equipment. Such local resources facilitate the satisfying of constraints placed on the system. For example, the local compute and storage resources allow an edge system to perform computations in real-time or near real-time, which may be a consideration in low latency user-cases such as autonomous driving, video surveillance, and mobile media consumption. Additionally, these resources will benefit from service management in an edge system which provides the ability to scale and achieve local service level agreements (SLAs) or service level objectives (SLOs), manage tiered service requirements, and enable local features and functions on a temporary or permanent basis.

A pool can include a device on a same chassis or different physically dispersed devices on different chassis or different racks. A resource pool can include homogeneous processors, homogeneous processors, and/or a memory pool.

An illustrative edge computing system may support and/or provide various services to endpoint devices (e.g., client user equipment (UEs)), one or more of which may have different requirements or constraints. For example, some services may have priority or quality-of-service (QoS) constraints (e.g., traffic data for autonomous vehicles may have a higher priority than temperature sensor data), reliability and resiliency (e.g., traffic data may require mission-critical reliability, while temperature data may be allowed some error variance), as well as power, cooling, and form-factor constraints. These and other technical constraints may offer significant complexity and technical challenges when applied in the multi-stakeholder setting.

FIG. **10** generically depicts an edge computing system **1000** for providing edge services and applications to multi-stakeholder entities, as distributed among one or more client compute nodes **1002**, one or more edge gateway nodes **1012**, one or more edge aggregation nodes **1022**, one or more core data centers **1032**, and a global network cloud **1042**, as distributed across layers of the network. The implementation of the edge computing system **1000** may be provided at or on behalf of a telecommunication service provider (“telco”, or “TSP”), internet-of-things service provider, cloud service provider (CSP), enterprise entity, or any other number of entities. Various implementations and configurations of the system **1000** may be provided dynamically, such as when orchestrated to meet service objectives.

For example, the client compute nodes **1002** are located at an endpoint layer, while the edge gateway nodes **1012** are located at an edge devices layer (local level) of the edge computing system **1000**. Additionally, the edge aggregation nodes **1022** (and/or fog devices **1024**, if arranged or operated with or among a fog networking configuration **1026**) are

located at a network access layer (an intermediate level). Fog computing (or “fogging”) generally refers to extensions of cloud computing to the edge of an enterprise’s network or to the ability to manage transactions across the cloud/edge landscape, typically in a coordinated distributed or multi-node network. Some forms of fog computing provide the deployment of compute, storage, and networking services between end devices and cloud computing data centers, on behalf of the cloud computing locations. Some forms of fog computing also provide the ability to manage the workload/workflow level services, in terms of the overall transaction, by pushing certain workloads to the edge or to the cloud based on the ability to fulfill the overall service level agreement. Fog computing in many scenarios provide a decentralized architecture and serves as an extension to cloud computing by collaborating with one or more edge node devices, providing the subsequent amount of localized control, configuration and management, and much more for end devices. Thus, some forms of fog computing provide operations that are consistent with edge computing as discussed herein; the edge computing aspects discussed herein are also applicable to fog networks, fogging, and fog configurations. Further, aspects of the edge computing systems discussed herein may be configured as a fog, or aspects of a fog may be integrated into an edge computing architecture.

The core data center **1032** is located at a core network layer (a regional or geographically-central level), while the global network cloud **1042** is located at a cloud data center layer (a national or world-wide layer). The use of “core” is provided as a term for a centralized network location-deeper in the network-which is accessible by multiple edge nodes or components; however, a “core” does not necessarily designate the “center” or the deepest location of the network. Accordingly, the core data center **1032** may be located within, at, or near the edge cloud **1000**. Although an illustrative number of client compute nodes **1002**, edge gateway nodes **1012**, edge aggregation nodes **1022**, edge core data centers **1032**, global network clouds **1042** are shown in FIG. **10**, it should be appreciated that the edge computing system **1000** may include additional devices or systems at one or more layers. Devices at a layer can be configured as peer nodes to another and, accordingly, act in a collaborative manner to meet service objectives.

Consistent with the examples provided herein, a client compute node **1002** may be embodied as any type of endpoint component, device, appliance, or other thing capable of communicating as a producer or consumer of data. Further, the label “node” or “device” as used in the edge computing system **1000** does not necessarily mean that such node or device operates in a client or agent/minion/follower role; rather, one or more of the nodes or devices in the edge computing system **1000** refer to individual entities, nodes, or subsystems which include discrete or connected hardware or software configurations to facilitate or use the edge cloud **1000**.

As such, the edge cloud **1000** is formed from network components and functional features operated by and within the edge gateway nodes **1012** and the edge aggregation nodes **1022**. The edge cloud **1000** may be embodied as any type of network that provides edge computing and/or storage resources which are proximately located to radio access network (RAN) capable endpoint devices (e.g., mobile computing devices, IoT devices, smart devices, etc.), which are shown in FIG. **10** as the client compute nodes **1002**. In other words, the edge cloud **1000** may be envisioned as an “edge” which connects the endpoint devices and traditional network access points that serves as an ingress point into service

provider core networks, including mobile carrier networks (e.g., Global System for Mobile Communications (GSM) networks, Long-Term Evolution (LTE) networks, 5G/6G networks, etc.), while also providing storage and/or compute capabilities. Other types and forms of network access (e.g., Wi-Fi, long-range wireless, wired networks including optical networks) may also be utilized in place of or in combination with such 3GPP carrier networks.

In some examples, the edge cloud **1000** may form a portion of or otherwise provide an ingress point into or across a fog networking configuration **1026** (e.g., a network of fog devices **1024**, not shown in detail), which may be embodied as a system-level horizontal and distributed architecture that distributes resources and services to perform a specific function. For instance, a coordinated and distributed network of fog devices **1024** may perform computing, storage, control, or networking aspects in the context of an IoT system arrangement. Other networked, aggregated, and distributed functions may exist in the edge cloud **1000** between the core data center **1032** and the client endpoints (e.g., client compute nodes **1002**). Some of these are discussed in the following sections in the context of network functions or service virtualization, including the use of virtual edges and virtual services which are orchestrated for multiple stakeholders.

As discussed in more detail below, the edge gateway nodes **1012** and the edge aggregation nodes **1022** cooperate to provide various edge services and security to the client compute nodes **1002**. Furthermore, because a client compute node **1002** may be stationary or mobile, a respective edge gateway node **1012** may cooperate with other edge gateway devices to propagate presently provided edge services, relevant service data, and security as the corresponding client compute node **1002** moves about a region. To do so, the edge gateway nodes **1012** and/or edge aggregation nodes **1022** may support multiple tenancy and multiple stakeholder configurations, in which services from (or hosted for) multiple service providers, owners, and multiple consumers may be supported and coordinated across a single or multiple compute devices.

A variety of security approaches may be utilized within the architecture of the edge cloud **1000**. In a multi-stakeholder environment, there can be multiple loadable security modules (LSMs) used to provision policies that enforce the stakeholder’s interests. Enforcement point environments could support multiple LSMs that apply the combination of loaded LSM policies (e.g., where the most constrained effective policy is applied, such as where if one or more of A, B or C stakeholders restricts access then access is restricted). Within the edge cloud **1000**, one or more edge entities can provision LSMs that enforce the Edge entity interests. The Cloud entity can provision LSMs that enforce the cloud entity interests. Likewise, the various Fog and IoT network entities can provision LSMs that enforce the Fog entity’s interests.

FIG. **11** shows an example where various client endpoints **1110** (in the form of mobile devices, computers, autonomous vehicles, business computing equipment, industrial processing equipment) provide requests **1120** for services or data transactions, and receive responses **1130** for the services or data transactions, to and from the edge cloud **1100** (e.g., via a wireless or wired network **1140**). Within the edge cloud **1000**, the CSP may deploy various compute and storage resources, such as edge content nodes **1150** to provide cached content from a distributed content delivery network. Other available compute and storage resources available on the edge content nodes **1150** may be used to execute other

15

services and fulfill other workloads. The edge content nodes **1150** and other systems of the edge cloud **1000** are connected to a cloud or data center **1170**, which uses a backhaul network **1160** to fulfill higher-latency requests from a cloud/data center for websites, applications, database servers, etc.

Examples described herein can be used in one or more examples of FIGS. 1-11 to allocate resources to perform one or more sections, stages or layers of a workload based on a complexity estimation from a pre-processing of data associated with the workload.

Pre-Processing Data to Determine Resources

FIG. 12 depicts an example edge cloud base station system. Multiple virtual network functions (VNFs) and services execute on various devices and pooled accelerators. Thus, through dynamic platform resource compositions, such as composite nodes, services can have access to diverse accelerators with different compute and quality of service (QoS) properties. In this context, artificial intelligence (AI) training and inference processing can be performed by accelerator devices at virtual edge appliances. In some cases, orchestration and resource management to perform services in accordance with QoS using software may be limited in scaling up if a number of services increases, and may not achieve real time performance requirements according to applicable QoS. In other words, as a number and variety of available accelerators and compute options increases, management of performance of services when an array of available accelerators and compute options is available may not achieve desired time-to-completion. One or more examples herein can be used in a communications base station such as but not limited to a 5G base station.

A platform architecture can be configured to generate meta-data indicative of processing resources to allocate for subsequent data processing to achieve applicable QoS. A workload request can trigger execution of one or more scheduling policies to select one or more accelerators and other device resources to allocate to perform a particular stage of a multiple level processing of data based on a QoS or service level agreement (SLA). A level of complexity to process data in a subsequent level or layer and output meta-data can be determined or estimated. Meta-data can indicate an estimated level of complexity or amount of computation to perform a subsequent level or layer of processing. For example, a neural network can have different layers that utilize different hardware resources to complete a layer of processing and pre-processing data of a workload request can identify various hardware resource configuration that could be used to perform the layer of processing within the applicable QoS or SLA. For a particular layer of processing, resource consumption can depend on a level of computation to be performed and pre-processing the data can identify a level of computation to allocate to perform processing of the layer. Latency of a completion of a layer can depend on the selected accelerator(s) or computing devices and load on the selected accelerator(s) or computing devices.

For example, examples can be used for artificial intelligence (AI) inference solutions where the complexity of phase N+1 typically depends on processing at phase N. Various examples, of AI inference include, but are not limited to: surveillance operations where person detection occurs at phase N and person identification occurs at phase N+1; smart cities where car detection occurs at phase N and license plate identification occurs at phase N+1; and speech analytics where speech to text occurs at phase N and natural language processing (NLP) occurs at phase N+1.

16

FIG. 13 depicts an example system. During, before, or after a processing stage of one or more processing stages, complexity of a subsequent processing stage can be estimated and determined and the complexity determination can be provided to a scheduler to determine hardware resources to utilize to perform the subsequent processing stage. Pre-processing can be performed for an AI or machine learning (ML) neural network that utilizes multiple layers and before scheduling hardware resources to perform processing associated with a layer.

A work request received at a platform or originated at the platform may identify an AI model type associated to be performed for the request and an SLA. The SLA may specify one or more of: AI model type, time to complete a pre-processing model, time to complete processing and make data available after a particular processing layer or stage, minimum compute resources, minimum network interface bandwidth to be allocated, minimum compute resources, minimum memory read or write rates. Examples of types of target operations include, but are not limited, to: person detection, object detection, music detection, and so forth. In response to receipt of a request to perform AI or ML inference, complexity estimator **1302** can identify a pre-processing model to be executed for the operation targeted by the request. For example, if the request indicates an AI model type includes identification of a person, complexity estimator **1302** could identify an associated pre-processing model of person detection and provides inputs to the person detection model that include a pointer to a payload to process and an SLA parameter for the acceleration function.

By performing the pre-processing model, complexity estimator **1302** can identify hardware resources to perform a subsequent layer or stage to satisfy the provided SLA from various of hardware resources (e.g., FPGA accelerators, memory, network, CPU, XPU, GPU, IPU, DPU, and so forth) available to access from the platform. Complexity estimator **1302** can output meta-data that indicates a complexity level for one or more layers performed for the request. A complexity level can include target priority and estimated runtime **1304** and estimated resources **1306**. For example, target priority and estimated runtime **1304** can indicate one or multiple devices and a time to completion for a subsequent phase. For example, target priority and estimated runtime **1304** can indicate accelerator device Accel0 can complete processing a subsequent phase in 10 milliseconds (ms), accelerator device Accel1 can complete processing a subsequent phase in 15 ms, and a CPU can complete the subsequent phase in 100 ms. For example, estimated resources **1306** can indicate hardware resources to allocate to complete the subsequent phase in the specified time to complete processing. For example, Accel0 can utilize allocation of N number of accelerators and M memory bandwidth to complete processing a subsequent phase in 10 ms; Accel1 can utilize allocation of P number of accelerators and Q memory bandwidth complete processing a subsequent phase in 15 ms; and use of a CPU can utilize R cores running at S frequency and T memory bandwidth complete processing a subsequent phase in 100 ms.

Load balancer **1308** may select from identified device resource allocations meta-data allocate the one or more hardware resources to perform the operations. Load balancer **1308** can implement one or more policies based on an applicable SLA to select one or more hardware resources identified in the meta-data to perform the target accelerated operation. Load balancer **1308** may receive telemetry from hardware resources and determine a level of resources available (e.g. memory or storage space availability,

memory bandwidth availability, compute capacity, number of available cores, accelerator resource availability, network device bandwidth availability, device temperatures and whether temperatures are near limits, etc.). Load balancer **1308** may estimate resources requirements to perform a next phase and decide which of the available hardware resources has a minimum available resources.

FIG. **14** depicts an example system. Load balancer **1404** to pre-process a workload and identify one or more combinations of hardware resources and configurations to perform one or more operations of layers of a multiple stage data processing to attempt to achieve applicable SLA parameters associated with the one or more layers. In some examples, the same or different implementations (e.g., different micro-service or software) of a layer can be executed by different devices. For example, load balancer **1404** can be implemented in one or more of: CPU, GPU, GPGPU, XPU, Quick Assist Technology (QAT), hardware queue manager (HQM), work scheduler, IPU, NIC, SmartNIC, software, container, virtual machine, or any other part of a platform.

Interface **1406** can be accessible from an operating system or driver and accessible by any software stack. Interface **1406** can permit a requester (e.g., device, orchestrator, application, container, virtual machine (VM), and so forth) to specify one or more of: (a) model identifier (ID); (b) model type; or (c) a bit-stream or binary associated with that model. For example, interface **1406** can include an application program interface (API), command line interface (CLI), script, binary, kernel, and so forth. For instance, a model type could be person detection, object detection, speech-to-text, image recognition, and so forth. In some examples, the type could be a unique ID type used to match the request to a bitstream. The bitstream could be provided with the request or accessed from memory by a pointer. The bitstream can include data that defines a kernel that is executable by a device (e.g., accelerators **0** to **2**) to perform one or more operations. For example, a bit stream may include instructions for performing a cryptographic operation (e.g., encryption or decryption), an arithmetic operation, a hashing operation, and/or other operations. A requester can provide one or more of: (i) pointer to the payload associated with the request; (ii) associated SLA for the offloaded processing operation; (iii) model type associated with a pre-processing phase; and/or (iv) pointer to memory area where the result of the pre-processing is to be stored and where the load balancer or accelerator can fetch the meta-data generated by pre-processing. The transaction or request can be part of a service function chain (SFC) that includes a sequence of operations, or a microservice in some examples.

Models cache and table **1408** can store information provided with the request. In examples, models cache and table **1408** can register a model by identifying a model identifier (ID) and its executable bitstream. Different binaries can execute on different accelerator devices in some examples and accordingly multiple binaries can be associated with a particular model type.

Interface **1406** can provide transactions or requests from one or more requesters into a queue among queues **0** to **X** to store a sequence of jobs to execute. In some examples, a queue can be associated with one or multiple accelerators among accelerators **1450-0** to **1450-2**. Although three accelerators are shown, any number of accelerators or other devices can be used.

Accelerator to perform complexity determination **1410** can execute a binary of a pre-processing model. Accelerator to perform complexity determination **1410** can be implemented as any type of processor described herein. For

example, a binary can indicate a complexity to measure and outputs a complexity factor. Complexity factor can to indicate hardware resource allocations needed to satisfy an SLA for different types of hardware devices such as execution on GPU, FPGA, CPU, and so forth.

Inputs to the pre-processing model can include the payload to be processed, the relevant SLA parameters and a pointer to a memory region meta-data is to be stored. Meta-data can indicate one or more configuration of resources to satisfy the provided SLA based on the various devices and accelerators connected to the platform. For example, meta-data can include a complexity factor that indicates a configuration of compute resources that could be used to complete a work based on an applicable SLA. For example, for a data set that is identified by pre-processing to not be sparse, a next layer can be executed on CPU to satisfy an SLA whereas a data set that is identified by pre-processing to be sparse, a next layer can be executed on an accelerator (e.g., FPGA).

Accelerator **1410** can enqueue a request after pre-processing the request and its data to determine a complexity of a job to a corresponding queue. Meta-data can be enqueued in a queue among queues **0** to **X** for access by one or more hardware resources. A set of queues can be devoted to store already pre-processed request for the various types of models (e.g., AI models). In some examples, a single queue can be used to store work requests for different AI models and/or multiple queues can be used where a queue is associated with performance of a particular AI model type. For example, a queues can be associated with a particular type of target operation.

Work requests can be in stored in a queue in order or out of order of execution by a device. A request can have an associated time stamp when it was inserted in queue and an associated SLA parameter(s) and the resource requirement meta-data. Scheduler **1412** can schedule performance of requests in order to try to achieve associated SLA parameter(s).

An accelerator among accelerators **1450-0** to **1450-2** can provide an input of a current amount of available resources to scheduler **1412** and type of accelerated function that it can execute and retrieve or pull content from one or more queues among queues **0** to **X**. For example, for accelerators **1450-0** to **1450-2**, respective pull and QoS circuitry **1452-0** to **1452-2** can be used to provide an input of a current amount of available resources to scheduler **1412** and type of accelerated function that it can execute and retrieve or pull content from one or more queues among queues **0** to **X**. An accelerator among accelerators **1450-0** to **1450-2** can retrieve or pull the meta-data generated by accelerator **1410** to determine an amount of resources to perform satisfy applicable QoS and SLAs. In some examples, pull and QoS circuitry of an accelerator can allocate accelerator resources, memory resources, and/or networking bandwidth to perform a workload.

In some examples, scheduler **1412** can process requests to retrieve meta-data and work requests from one or more of accelerators **1450-0** to **1450-2** and decide which device is to execute a model based on QoS and/or SLA and select a device to execute model based on complexity (e.g., compute resources required to complete work to perform in required time, number of uncompleted jobs, processing speed (e.g., flops/second), and/or time to complete a particular layer according to QoS or SLA. In some examples, different combinations of resource requirements can be specified for an accelerator to achieve different levels of an SLA and load balancer may determine to delay performance of a request

(e.g., by 100 ns) at the expense utilizing 10% more of accelerator resource to satisfy an applicable SLA parameter. In some cases, scheduler **1412** can register or de-register models depending respective available hardware resources or lack of available hardware resources.

FIG. **15** depicts an example operation. In a surveillance and person identification AI model, pre-processing NN can determine complexity of surveillance and person identification to be performed in layer 0. For example, pre-processing NN can determine a number of persons in the image. In this example, Resource A is used to perform pre-processing NN and can utilize 2 GB/s of DDR memory and 2 vision processing unit (VPU) cores). Pre-processing can generate different combinations of resource requirements for the same accelerators providing different levels of SLA. For example, Resources B-D can be identified where Resource B includes 2 Gbps DDR memory and 2 VPU cores, Resource C includes 0.5 Gbps DDR memory and 1 VPU core, and Resource D includes 3.5 Gbps DDR memory and 5 VPU cores.

Layer 0 of the NN can identify the specific individuals that may match one or more of the identified persons and Layer 1 and Layer X can perform other operations. In this example, Layer 0 of NN has different level of compute requirements depending on the meta-data generated in the pre-processing NN layer. Hence, resources can be allocated prior to execution of the first and/or second phase in order to meet the required SLA (e.g., 40 fps). In some cases, the SLA may change over time for the stream depending external facts such as camera location, flow of people being analyzed, etc., and resources allocated to achieve the SLA can be adjusted accordingly.

FIG. **16** depicts an example process. The process can be performed by a device that determines processing to determine one or more configurations of device resources to select from to perform a subsequent layer of processing. At **1602**, a request can be received at a pre-processing stage to perform processing of data. In some examples, the request includes a binary that is performed in a pre-processing stage and the request can include an SLA. At **1604**, the pre-processing stage can estimate or determine device resources to complete a next stage within a requirement of an applicable SLA. For example, one or multiple configurations of device resources can be provided by the pre-processing stage. The one or multiple configurations of device resources can be indicated in meta-data.

At **1606**, a workload associated with the request can be provided to a queue with pre-processing meta-data. The compute resources associated with the workload can be indicated in the meta-data. In some examples, a hardware device can allocate resources to perform the workload. A queue can be associated with a particular AI model type in some examples. At **1608**, a device can read the queue to identify a workload to perform. The device can be configured to allocate hardware resources specified in the meta-data and perform the workload. Performing the workload can generate data that can be processed by a subsequent layer or made available to the requester.

FIG. **17** depicts an example computing system. System **1700** can determine one or more configurations of hardware to perform a stage or a workload in accordance with one or more embodiments described herein. System **1700** includes processor **1710**, which provides processing, operation management, and execution of instructions for system **1700**. Processor **1710** can include any type of microprocessor, central processing unit (CPU), graphics processing unit (GPU), vision processing unit (VPU), processing core, or

other processing hardware to provide processing for system **1700**, or a combination of processors. Processor **1710** controls the overall operation of system **1700**, and can be or include, one or more programmable general-purpose or special-purpose microprocessors, digital signal processors (DSPs), programmable controllers, application specific integrated circuits (ASICs), programmable logic devices (PLDs), or the like, or a combination of such devices.

In one example, system **1700** includes interface **1712** coupled to processor **1710**, which can represent a higher speed interface or a high throughput interface for system components that needs higher bandwidth connections, such as memory subsystem **1720** or graphics interface components **1740**, or accelerators **1742**. Interface **1712** represents an interface circuit, which can be a standalone component or integrated onto a processor die. Where present, graphics interface **1740** interfaces to graphics components for providing a visual display to a user of system **1700**. In one example, graphics interface **1740** can drive a high definition (HD) display that provides an output to a user. High definition can refer to a display having a pixel density of approximately 100 PPI (pixels per inch) or greater and can include formats such as full HD (e.g., 1080p), retina displays, 4K (ultra-high definition or UHD), or others. In one example, the display can include a touchscreen display. In one example, graphics interface **1740** generates a display based on data stored in memory **1730** or based on operations executed by processor **1710** or both. In one example, graphics interface **1740** generates a display based on data stored in memory **1730** or based on operations executed by processor **1710** or both.

Accelerators **1742** can be a fixed function or programmable offload engine that can be accessed or used by a processor **1710**. For example, an accelerator among accelerators **1742** can provide compression (DC) capability, cryptography services such as public key encryption (PKE), cipher, hash/authentication capabilities, decryption, or other capabilities or services. In addition or alternatively, an accelerator among accelerators **1742** provides field select controller capabilities as described herein. In some cases, accelerators **1742** can be integrated into a CPU socket (e.g., a connector to a motherboard or circuit board that includes a CPU and provides an electrical interface with the CPU). For example, accelerators **1742** can include a single or multi-core processor, graphics processing unit, logical execution unit single or multi-level cache, functional units usable to independently execute programs or threads, application specific integrated circuits (ASICs), neural network processors (NNPs), programmable control logic, and programmable processing elements such as field programmable gate arrays (FPGAs) or programmable logic devices (PLDs). Accelerators **1742** can provide multiple neural networks, CPUs, processor cores, general purpose graphics processing units, or graphics processing units can be made available for use by artificial intelligence (AI) or machine learning (ML) models. For example, the AI model can use or include one or more of: a reinforcement learning scheme, Q-learning scheme, deep-Q learning, or Asynchronous Advantage Actor-Critic (A3C), combinatorial neural network, recurrent combinatorial neural network, or other AI or ML model. Multiple neural networks, processor cores, or graphics processing units can be made available for use by AI or ML models.

Memory subsystem **1720** represents the main memory of system **1700** and provides storage for code to be executed by processor **1710**, or data values to be used in executing a routine. Memory subsystem **1720** can include one or more

memory devices **1730** such as read-only memory (ROM), flash memory, one or more varieties of random access memory (RAM) such as DRAM, or other memory devices, or a combination of such devices. Memory **1730** stores and hosts, among other things, operating system (OS) **1732** to provide a software platform for execution of instructions in system **1700**. Additionally, applications **1734** can execute on the software platform of OS **1732** from memory **1730**. Applications **1734** represent programs that have their own operational logic to perform execution of one or more functions. Processes **1736** represent agents or routines that provide auxiliary functions to OS **1732** or one or more applications **1734** or a combination. OS **1732**, applications **1734**, and processes **1736** provide software logic to provide functions for system **1700**. In one example, memory subsystem **1720** includes memory controller **1722**, which is a memory controller to generate and issue commands to memory **1730**. It will be understood that memory controller **1722** could be a physical part of processor **1710** or a physical part of interface **1712**. For example, memory controller **1722** can be an integrated memory controller, integrated onto a circuit with processor **1710**.

In some examples, OS **1732** can be Linux®, Windows® Server or personal computer, FreeBSD®, Android®, MacOS®, iOS®, VMware vSphere, openSUSE, RHEL, CentOS, Debian, Ubuntu, or any other operating system. The OS and driver can execute on a CPU sold or designed by Intel®, ARM®, AMD®, Qualcomm®, IBM®, Texas Instruments®, among others.

While not specifically illustrated, it will be understood that system **1700** can include one or more buses or bus systems between devices, such as a memory bus, a graphics bus, interface buses, or others. Buses or other signal lines can communicatively or electrically couple components together, or both communicatively and electrically couple the components. Buses can include physical communication lines, point-to-point connections, bridges, adapters, controllers, or other circuitry or a combination. Buses can include, for example, one or more of a system bus, a Peripheral Component Interconnect (PCI) bus, a Hyper Transport or industry standard architecture (ISA) bus, a small computer system interface (SCSI) bus, a universal serial bus (USB), or an Institute of Electrical and Electronics Engineers (IEEE) standard 1394 bus (Firewire).

In one example, system **1700** includes interface **1714**, which can be coupled to interface **1712**. In one example, interface **1714** represents an interface circuit, which can include standalone components and integrated circuitry. In one example, multiple user interface components or peripheral components, or both, couple to interface **1714**. Network interface **1750** provides system **1700** the ability to communicate with remote devices (e.g., servers or other computing devices) over one or more networks. Network interface **1750** can include an Ethernet adapter, wireless interconnection components, cellular network interconnection components, USB (universal serial bus), or other wired or wireless standards-based or proprietary interfaces. Network interface **1750** can transmit data to a device that is in the same data center or rack or a remote device, which can include sending data stored in memory. Network interface **1750** can receive data from a remote device, which can include storing received data into memory.

In one example, system **1700** includes one or more input/output (I/O) interface(s) **1760**. I/O interface **1760** can include one or more interface components through which a user interacts with system **1700** (e.g., audio, alphanumeric, tactile/touch, or other interfacing). Peripheral interface **1770**

can include any hardware interface not specifically mentioned above. Peripherals refer generally to devices that connect dependently to system **1700**. A dependent connection is one where system **1700** provides the software platform or hardware platform or both on which operation executes, and with which a user interacts.

In one example, system **1700** includes storage subsystem **1780** to store data in a nonvolatile manner. In one example, in certain system implementations, at least certain components of storage **1780** can overlap with components of memory subsystem **1720**. Storage subsystem **1780** includes storage device(s) **1784**, which can be or include any conventional medium for storing large amounts of data in a nonvolatile manner, such as one or more magnetic, solid state, or optical based disks, or a combination. Storage **1784** holds code or instructions and data **1786** in a persistent state (e.g., the value is retained despite interruption of power to system **1700**). Storage **1784** can be generically considered to be a “memory,” although memory **1730** is typically the executing or operating memory to provide instructions to processor **1710**. Whereas storage **1784** is nonvolatile, memory **1730** can include volatile memory (e.g., the value or state of the data is indeterminate if power is interrupted to system **1700**). In one example, storage subsystem **1780** includes controller **1782** to interface with storage **1784**. In one example controller **1782** is a physical part of interface **1714** or processor **1710** or can include circuits or logic in both processor **1710** and interface **1714**.

A volatile memory is memory whose state (and therefore the data stored in it) is indeterminate if power is interrupted to the device. Dynamic volatile memory uses refreshing the data stored in the device to maintain state. One example of dynamic volatile memory includes DRAM (Dynamic Random Access Memory), or some variant such as Synchronous DRAM (SDRAM). An example of a volatile memory includes a cache. A memory subsystem as described herein may be compatible with a number of memory technologies, such as DDR3 (Double Data Rate version 3, original release by JEDEC (Joint Electronic Device Engineering Council) on Jun. 16, 2007). DDR4 (DDR version 4, initial specification published in September 2012 by JEDEC), DDR4E (DDR version 4), LPDDR3 (Low Power DDR version 3, JESD209-3B, August 2013 by JEDEC), LPDDR4 (LPDDR version 4, JESD209-4, originally published by JEDEC in August 2014), WIO2 (Wide Input/output version 2, JESD229-2 originally published by JEDEC in August 2014, HBM (High Bandwidth Memory, JESD325, originally published by JEDEC in October 2013, LPDDR5 (currently in discussion by JEDEC), HBM2 (HBM version 2), currently in discussion by JEDEC, or others or combinations of memory technologies, and technologies based on derivatives or extensions of such specifications.

A non-volatile memory (NVM) device is a memory whose state is determinate even if power is interrupted to the device. The NVM device can comprise a block addressable memory device, such as NAND technologies, or more specifically, multi-threshold level NAND flash memory (for example, Single-Level Cell (“SLC”), Multi-Level Cell (“MLC”), Quad-Level Cell (“QLC”), Tri-Level Cell (“TLC”), or some other NAND). A NVM device can also comprise a byte-addressable write-in-place three dimensional cross point memory device, or other byte addressable write-in-place NVM device (also referred to as persistent memory), such as single or multi-level Phase Change Memory (PCM) or phase change memory with a switch (PCMS), Intel® Optane™ memory, NVM devices that use chalcogenide phase change material (for example, chalco-

genide glass), resistive memory including metal oxide base, oxygen vacancy base and Conductive Bridge Random Access Memory (CB-RAM), nanowire memory, ferroelectric random access memory (FeRAM, FRAM), magnetoresistive random access memory (MRAM) that incorporates memristor technology, spin transfer torque (STT)-MRAM, a spintronic magnetic junction memory based device, a magnetic tunneling junction (MTJ) based device, a DW (Domain Wall) and SOT (Spin Orbit Transfer) based device, a thyristor based memory device, or a combination of one or more of the above, or other memory.

A power source (not depicted) provides power to the components of system 1700. More specifically, power source typically interfaces to one or multiple power supplies in system 1700 to provide power to the components of system 1700. In one example, the power supply includes an AC to DC (alternating current to direct current) adapter to plug into a wall outlet. Such AC power can be renewable energy (e.g., solar power) power source. In one example, power source includes a DC power source, such as an external AC to DC converter. In one example, power source or power supply includes wireless charging hardware to charge via proximity to a charging field. In one example, power source can include an internal battery, alternating current supply, motion-based power supply, solar power supply, or fuel cell source.

In an example, system 1700 can be implemented using interconnected compute sleds of processors, memories, storages, network interfaces, and other components. High speed interconnects can be used such as: Ethernet (IEEE 802.3), remote direct memory access (RDMA), InfiniBand, Internet Wide Area RDMA Protocol (iWARP), Transmission Control Protocol (TCP), User Datagram Protocol (UDP), quick UDP Internet Connections (QUIC), RDMA over Converged Ethernet (RoCE), Peripheral Component Interconnect express (PCIe), Intel QuickPath Interconnect (QPI), Intel Ultra Path Interconnect (UPI), Intel On-Chip System Fabric (IOSF), Omni-Path, Compute Express Link (CXL), HyperTransport, high-speed fabric, NVLink, Advanced Microcontroller Bus Architecture (AMBA) interconnect, OpenCAPI, Gen-Z, Infinity Fabric (IF), Cache Coherent Interconnect for Accelerators (CCIX), 3GPP Long Term Evolution (LTE) (4G), 3GPP 5G, and variations thereof. Data can be copied or stored to virtualized storage nodes or accessed using a protocol such as NVMe over Fabrics (NVMe-oF) or NVMe.

Examples herein can be part of computing and networking equipment, such as switches, routers, racks, and blade servers such as those employed in a data center and/or server farm environment. The servers used in data centers and server farms comprise arrayed server configurations such as rack-based servers or blade servers. These servers are interconnected in communication via various network provisions, such as partitioning sets of servers into Local Area Networks (LANs) with appropriate switching and routing facilities between the LANs to form a private Intranet. For example, cloud hosting facilities may typically employ large data centers with a multitude of servers. A blade comprises a separate computing platform that is configured to perform server-type functions, that is, a “server on a card.” Accordingly, a blade can include components common to conventional servers, including a main printed circuit board (main board) providing internal wiring (e.g., buses) for coupling appropriate integrated circuits (ICs) and other components mounted to the board.

In some examples, network interface and other examples described herein can be used in connection with a base station (e.g., 3G, 4G, 5G and so forth), macro base station

(e.g., 5G networks), picostation (e.g., an IEEE 802.11 compatible access point), nanostation (e.g., for Point-to-Multi-Point (PtMP) applications), on-premises data centers, off-premises data centers, edge network elements, edge servers, edge switches, fog network elements, and/or hybrid data centers (e.g., data center that use virtualization, cloud and software-defined networking to deliver application workloads across physical data centers and distributed multi-cloud environments).

Various examples may be implemented using hardware elements, software elements, or a combination of both. In some examples, hardware elements may include devices, components, processors, microprocessors, circuits, circuit elements (e.g., transistors, resistors, capacitors, inductors, and so forth), integrated circuits, ASICs, PLDs, DSPs, FPGAs, memory units, logic gates, registers, semiconductor device, chips, microchips, chip sets, and so forth. In some examples, software elements may include software components, programs, applications, computer programs, application programs, system programs, machine programs, operating system software, middleware, firmware, software modules, routines, subroutines, functions, methods, procedures, software interfaces, APIs, instruction sets, computing code, computer code, code segments, computer code segments, words, values, symbols, or combination thereof. Determining whether an example is implemented using hardware elements and/or software elements may vary in accordance with any number of factors, such as desired computational rate, power levels, heat tolerances, processing cycle budget, input data rates, output data rates, memory resources, data bus speeds and other design or performance constraints, as desired for a given implementation. A processor can be one or more combination of a hardware state machine, digital control logic, central processing unit, or any hardware, firmware and/or software elements.

Some examples may be implemented using or as an article of manufacture or at least one computer-readable medium. A computer-readable medium may include a non-transitory storage medium to store logic. In some examples, the non-transitory storage medium may include one or more types of computer-readable storage media capable of storing electronic data, including volatile memory or non-volatile memory, removable or non-removable memory, erasable or non-erasable memory, writeable or re-writable memory, and so forth. In some examples, the logic may include various software elements, such as software components, programs, applications, computer programs, application programs, system programs, machine programs, operating system software, middleware, firmware, software modules, routines, subroutines, functions, methods, procedures, software interfaces, API, instruction sets, computing code, computer code, code segments, computer code segments, words, values, symbols, or combination thereof.

According to some examples, a computer-readable medium may include a non-transitory storage medium to store or maintain instructions that when executed by a machine, computing device or system, cause the machine, computing device or system to perform methods and/or operations in accordance with the described examples. The instructions may include any suitable type of code, such as source code, compiled code, interpreted code, executable code, static code, dynamic code, and the like. The instructions may be implemented according to a predefined computer language, manner or syntax, for instructing a machine, computing device or system to perform a certain function. The instructions may be implemented using any suitable

high-level, low-level, object-oriented, visual, compiled and/or interpreted programming language.

One or more aspects of at least one example may be implemented by representative instructions stored on at least one machine-readable medium which represents various logic within the processor, which when read by a machine, computing device or system causes the machine, computing device or system to fabricate logic to perform the techniques described herein. Such representations, known as “IP cores” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

The appearances of the phrase “one example” or “an example” are not necessarily all referring to the same example or embodiment. Any aspect described herein can be combined with any other aspect or similar aspect described herein, regardless of whether the aspects are described with respect to the same figure or element. Division, omission or inclusion of block functions depicted in the accompanying figures does not infer that the hardware components, circuits, software and/or elements for implementing these functions would necessarily be divided, omitted, or included in embodiments.

Some examples may be described using the expression “coupled” and “connected” along with their derivatives. These terms are not necessarily intended as synonyms for each other. For example, descriptions using the terms “connected” and/or “coupled” may indicate that two or more elements are in direct physical or electrical contact with each other. The term “coupled,” however, may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

The terms “first,” “second,” and the like, herein do not denote any order, quantity, or importance, but rather are used to distinguish one element from another. The terms “a” and “an” herein do not denote a limitation of quantity, but rather denote the presence of at least one of the referenced items. The term “asserted” used herein with reference to a signal denote a state of the signal, in which the signal is active, and which can be achieved by applying any logic level either logic 0 or logic 1 to the signal. The terms “follow” or “after” can refer to immediately following or following after some other event or events. Other sequences of operations may also be performed according to alternatives. Furthermore, additional operations may be added or removed depending on the particular applications. Any combination of changes can be used and one of ordinary skill in the art with the benefit of this disclosure would understand the many variations, modifications, and alternatives thereof.

Disjunctive language such as the phrase “at least one of X, Y, or Z,” unless specifically stated otherwise, is otherwise understood within the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present. Additionally, conjunctive language such as the phrase “at least one of X, Y, and Z,” unless specifically stated otherwise, should also be understood to mean X, Y, Z, or combination thereof, including “X, Y, and/or Z.”

Illustrative examples of the devices, systems, and methods disclosed herein are provided below. An embodiment of the devices, systems, and methods may include one or more, and combination of, the examples described below.

Performing a first processing of data in a multi-level neural network (NN), where performing the first processing of data includes generating meta-data indicative of an amount of processing for a second processing of the data in the multi-level NN; selecting one or more hardware resources for performance of the second processing of the data; and performing, by the selected one or more hardware resources, the second processing of the data in the multi-level NN, wherein one or more hardware resources utilized by the second processing is different than one or more hardware resources utilized by the first processing.

Example 1 includes a method comprising: receiving a workload request to perform a pre-processing operation and at least one service level agreement (SLA) parameter and performing the pre-processing operation associated with the workload request and the at least one SLA parameter to determine multiple configurations of hardware resources to perform a workload associated with the workload request in a subsequent stage.

Example 2 includes one or more examples, wherein the pre-processing operation comprises an executable binary referenced by the workload request.

Example 3 includes one or more examples, and includes: storing, in a queue, the multiple configurations of hardware resources to perform a workload associated with the workload request in the subsequent stage, wherein the multiple configurations of hardware resources are available for access by one or more accelerator devices to perform the workload.

Example 4 includes one or more examples, and includes: selecting one or more hardware resources for performance of the workload associated with the workload request among the multiple configurations of hardware resources.

Example 5 includes one or more examples, wherein the subsequent stage comprises a data processing stage in a neural network (NN).

Example 6 includes one or more examples, wherein the workload associated with the workload request in the subsequent stage comprises a machine learning (ML) inference operation.

Example 7 includes one or more examples, and includes a non-transitory computer-readable medium comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: determine multiple configurations of hardware resources to perform a workload associated with a workload request in a subsequent stage based on a pre-processing operation associated with the workload request and at least one service level agreement (SLA) parameter associated with the workload request.

Example 8 includes one or more examples, wherein an executable binary is associated with the workload request and execution of the executable binary performs the pre-processing operation.

Example 9 includes one or more examples, and includes instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: store the multiple configurations of hardware resources to perform a workload associated with the workload request in a subsequent stage, wherein the multiple configurations of hardware resources are available for access by one or more accelerator devices to perform the workload.

Example 10 includes one or more examples, and includes instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: select one or more hardware resources for performance of the workload associated with the workload request based on one or more of the multiple configurations of hardware resources.

27

Example 11 includes one or more examples, wherein the subsequent stage comprises a data processing stage in a neural network (NN).

Example 12 includes one or more examples, wherein the workload associated with the workload request in a subsequent stage comprises a machine learning (ML) inference operation.

Example 13 includes one or more examples, and includes a system comprising: at least one processor and circuitry to: determine multiple configurations of hardware resources to perform a workload associated with a workload request in a subsequent stage based on a pre-processing operation associated with the workload request and at least one service level agreement (SLA) parameter associated with the workload request.

Example 14 includes one or more examples, wherein an executable binary is associated with the workload request and execution of the executable binary performs the pre-processing operation.

Example 15 includes one or more examples, wherein a processor of the at least one processor is to execute the executable binary.

Example 16 includes one or more examples, wherein the circuitry is to: store the multiple configurations of hardware resources to perform a workload associated with the workload request in a subsequent stage, wherein the multiple configurations of hardware resources are available for access by one or more accelerator devices to perform the workload.

Example 17 includes one or more examples, wherein the circuitry is to: select one or more hardware resources for performance of the workload associated with the workload request based on one or more of the multiple configurations of hardware resources.

Example 18 includes one or more examples, wherein the workload associated with the workload request in a subsequent stage comprises a machine learning (ML) inference operation.

Example 19 includes one or more examples, and includes a non-transitory computer-readable medium comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: receive a workload request to perform a pre-processing operation and at least one service level agreement (SLA) parameter and cause performance of the pre-processing operation associated with the workload request and the at least one SLA parameter to determine multiple configurations of hardware resources to perform a workload associated with the workload request in a subsequent stage.

Example 20 includes one or more examples, and includes instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: store the multiple configurations of hardware resources to perform a workload associated with the workload request in a subsequent stage, wherein the multiple configurations of hardware resources are available for access by one or more accelerator devices to perform the workload and select one or more hardware resources for performance of the workload associated with the workload request based on one or more of the multiple configurations of hardware resources.

The invention claimed is:

1. A method comprising:

receiving a request to perform a pre-processing operation and at least one service level agreement (SLA) parameter;

determining a first configuration of hardware resources to allocate to perform operations of a first layer of a neural network (NN) associated with the request based on the

28

at least one SLA parameter prior to scheduling and performance of the first layer of the NN by performing a pre-processing operation associated with the request; performing the operations of the first layer of the NN by use of the first configuration of the allocated hardware resources;

determining a second configuration of hardware resources to allocate to perform operations of a second layer of the NN associated with the request based on the at least one SLA parameter prior to scheduling and performance of the second layer of the NN by performing a second pre-processing operation associated with the request; and

performing the operations of the second layer of the NN by use of the second configuration of the allocated hardware resources.

2. The method of claim 1, wherein the pre-processing operation comprises an executable binary referenced by the request.

3. The method of claim 1, comprising:

storing, in a queue, the first and second configurations of hardware resources to perform operations associated with the request, wherein the first and second configurations of hardware resources are available for access by one or more accelerator devices for performing the operations.

4. The method of claim 1, comprising:

selecting one or more hardware resources for performing the operations associated with the request among the first and second configurations of hardware resources.

5. The method of claim 1, wherein the first layer of the NN comprises a data processing stage in the NN and wherein the NN comprises multiple stages.

6. The method of claim 1, wherein the operations associated with the request in the first layer of the NN comprise machine learning (ML) inference operations.

7. The method of claim 1, wherein the NN comprises the pre-processing operation.

8. At least one non-transitory computer-readable medium comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: prior to scheduling and performance of a first layer of a neural network (NN), perform a pre-processing operation to determine a first configuration of hardware resources to perform operations of the first layer of the NN associated with a request based on at least one service level agreement (SLA) parameter associated with the request;

perform the operations of the first layer of the NN by use of the first configuration of the hardware resources;

prior to scheduling and performance of a second layer of the NN, determine a second configuration of hardware resources to perform operations of the second layer of the NN associated with the request based on the at least one SLA parameter associated with the request; and perform the operations of the second layer of the NN by use of the second configuration of the hardware resources.

9. The computer-readable medium of claim 8, wherein an executable binary is associated with the request and wherein the performing the pre-processing operation comprises execution of the executable binary.

10. The computer-readable medium of claim 8, comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: store the first and second configurations of hardware resources to perform the operations associated with the

29

request, wherein the first and second configurations of hardware resources are available for access by one or more accelerator devices to perform the operations.

11. The computer-readable medium of claim 8, comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to:

select one or more hardware resources for performance of the operations associated with the request based on one or more of the first and second configurations of hardware resources.

12. The computer-readable medium of claim 8, wherein the first layer of the NN comprises a data processing stage in the NN and wherein the NN comprises multiple stages.

13. The computer-readable medium of claim 8, wherein the operations associated with the request in the first layer of the NN comprise machine learning (ML) inference operations.

14. A system comprising:

at least one processor and circuitry to:

prior to scheduling and performance of a first layer of a neural network (NN), perform a pre-processing operation to determine a first configuration of hardware resources to perform operations of the first layer of the NN associated with a request based on at least one service level agreement (SLA) parameter associated with the request;

perform the operations of the first layer of the NN by use of the first configuration of the hardware resources;

prior to scheduling and performance of a second layer of the NN, determine a second configuration of hardware resources to perform operations of the second layer of the NN associated with the request based on the at least one SLA parameter associated with the request; and

perform the operations of the second layer of the NN by use of the second configuration of the hardware resources.

15. The system of claim 14, wherein an executable binary is associated with the request and wherein the perform the pre-processing operation comprises execution of the executable binary.

16. The system of claim 15, wherein a processor of the at least one processor is to execute the executable binary.

17. The system of claim 14, wherein the circuitry is to: store the first and second configurations of hardware resources to perform the operations associated with the request, wherein the first and second configurations of

30

hardware resources are available for access by one or more accelerator devices to perform the operations.

18. The system of claim 14, wherein the circuitry is to: select one or more hardware resources for performance of the operations associated with the request based on one or more of the first and second configurations of hardware resources.

19. The system of claim 14, wherein the operations associated with the request in the first layer of the NN comprise machine learning (ML) inference operations.

20. At least one non-transitory computer-readable medium comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to:

receive a request to perform a pre-processing operation and at least one service level agreement (SLA) parameter;

prior to scheduling and performance of a first layer of a neural network (NN), determine a first configuration of hardware resources to perform operations of the first layer of the NN associated with the request in a subsequent stage, wherein the pre-processing operation is based on operations of a neural network (NN) and the at least one SLA parameter;

cause performance of the operations of the first layer of the NN by use of the first configuration of the hardware resources;

prior to scheduling and performance of a second layer of the NN, determine a second configuration of hardware resources to perform operations of the second layer of the NN associated with the request based on the at least one SLA parameter associated with the request; and perform the operations of the second layer of the NN by use of the second configuration of the hardware resources.

21. The computer-readable medium of claim 20, comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to:

store the first and second configurations of hardware resources to perform the operations associated with the request in the subsequent stage, wherein the first and second configurations of hardware resources are available for access by one or more accelerator devices to perform the operations and

select one or more hardware resources for performance of the operations associated with the request based on one or more of the first and second configurations of hardware resources.

* * * * *