



US012393504B2

(12) **United States Patent**
Deng et al.

(10) **Patent No.:** **US 12,393,504 B2**
(45) **Date of Patent:** ***Aug. 19, 2025**

(54) **AUTOMATED PROGRAM REPAIR TOOL**

(71) Applicant: **MICROSOFT TECHNOLOGY LICENSING, LLC.**, Redmond, WA (US)

(72) Inventors: **Shao Kun Deng**, Seattle, WA (US); **Neelakantan Sundaresan**, Bellevue, WA (US); **Alexey Svyatkovskiy**, Bellevue, WA (US); **Michele Tufano**, Bellevue, WA (US)

(73) Assignee: **Microsoft Technology Licensing, LLC**, Redmond, WA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 0 days.

This patent is subject to a terminal disclaimer.

(21) Appl. No.: **18/628,773**

(22) Filed: **Apr. 7, 2024**

(65) **Prior Publication Data**

US 2024/0370352 A1 Nov. 7, 2024

Related U.S. Application Data

(63) Continuation of application No. 17/994,185, filed on Nov. 25, 2022, now Pat. No. 11,977,474, which is a (Continued)

(51) **Int. Cl.**
G06F 11/362 (2025.01)
G06N 20/00 (2019.01)

(52) **U.S. Cl.**
CPC **G06F 11/362** (2013.01); **G06N 20/00** (2019.01)

(58) **Field of Classification Search**

CPC G06F 11/362; G06N 20/00
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

11,593,675 B1 * 2/2023 Garg G06F 21/54
2007/0022321 A1 * 1/2007 Chua G06F 11/0775
714/38.1

(Continued)

OTHER PUBLICATIONS

Md. Mostafizer Rahman et al. "Source Code Assessment and Classification Based on Estimated Error Probability Using Attentive LSTM Language Model and Its Application in Programming Education", [Online], pp. 1-21, [Retrieved from Internet on Apr. 18, 2025], <<https://www.mdpi.com/2076-3417/10/8/2973>> (Year: 2020).*

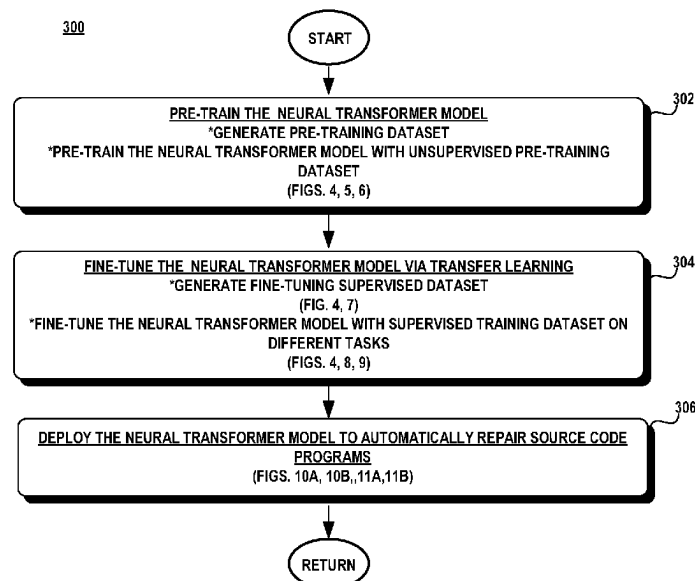
(Continued)

Primary Examiner — Ziaul A Chowdhury

(57) **ABSTRACT**

An automated program repair tool utilizes a neural transformer model with attention to predict the contents of a bug repair in the context of source code having a bug of an identified bug type. The neural transformer model is trained on a large unsupervised corpus of source code using a span-masking denoising optimization objective, and fine-tuned on a large supervised dataset of triplets containing a bug-type annotation, software bug, and repair. The bug-type annotation is derived from an interprocedural static code analyzer. A bug type edit centroid is computed for each bug type and used in the inference decoding phase to generate the bug repair.

17 Claims, 14 Drawing Sheets



Related U.S. Application Data

continuation of application No. 16/897,824, filed on
Jun. 10, 2020, now Pat. No. 11,526,424.

- (60) Provisional application No. 63/025,535, filed on May
15, 2020.

(56) **References Cited**

U.S. PATENT DOCUMENTS

2021/0182077 A1 * 6/2021 Chen G06Q 30/0241
2021/0311853 A1 * 10/2021 Shao G06F 11/3604

OTHER PUBLICATIONS

Communication pursuant to Article 94(3) received in European
Application No. 21720955.0, mailed on May 13, 2024, 12 pages.

* cited by examiner

100




FIG. 1

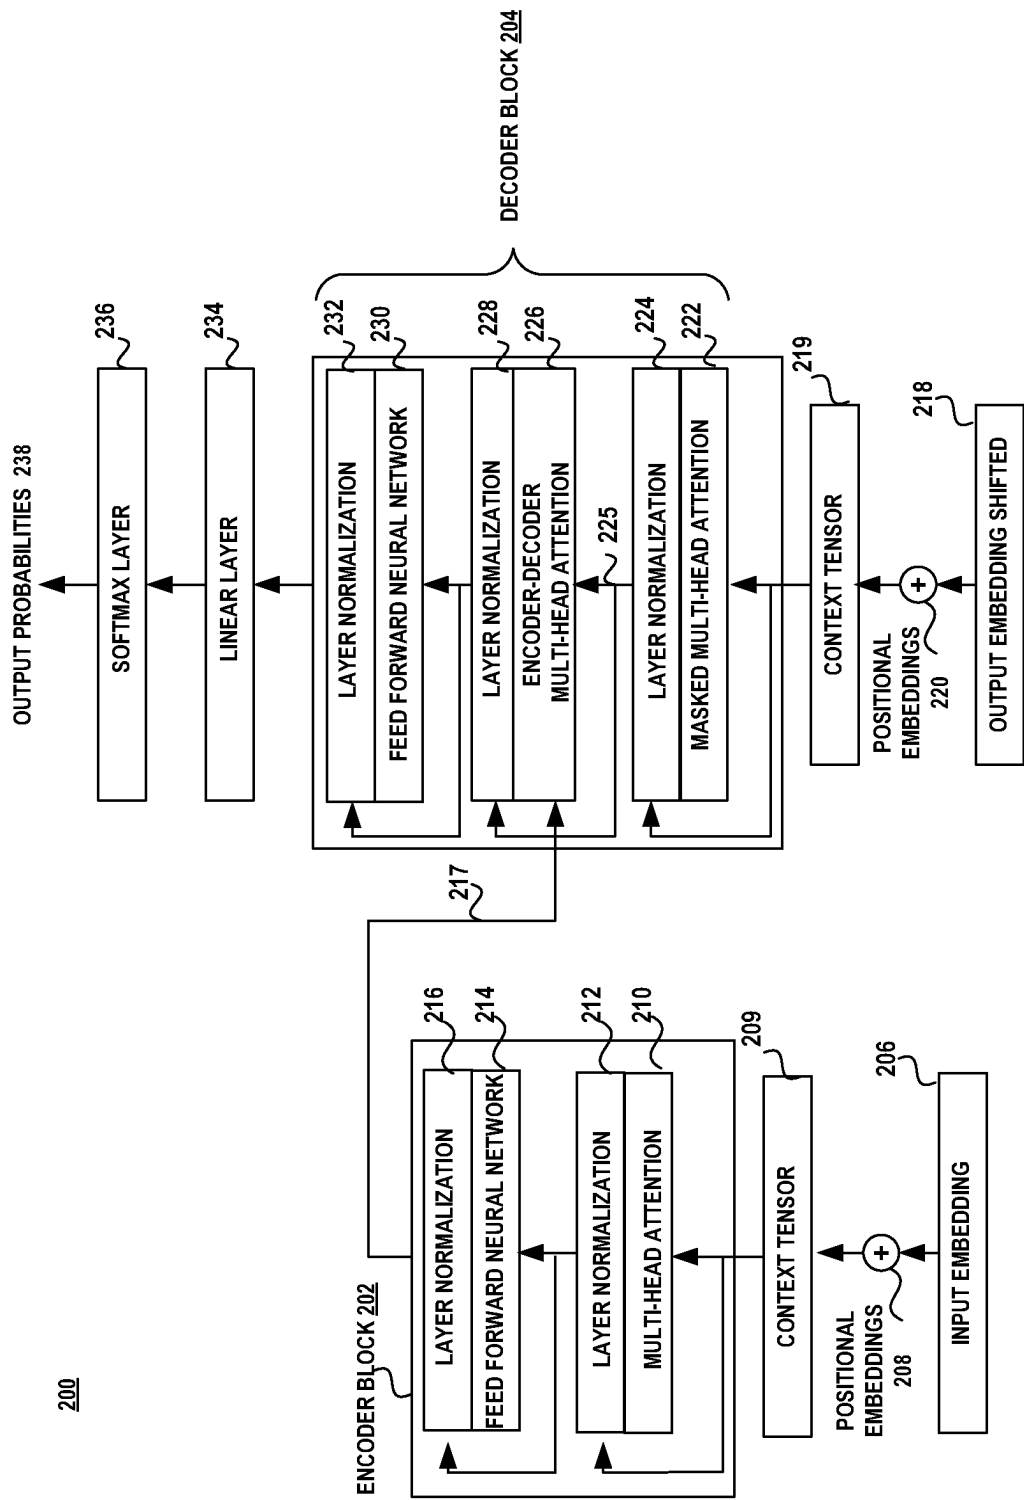
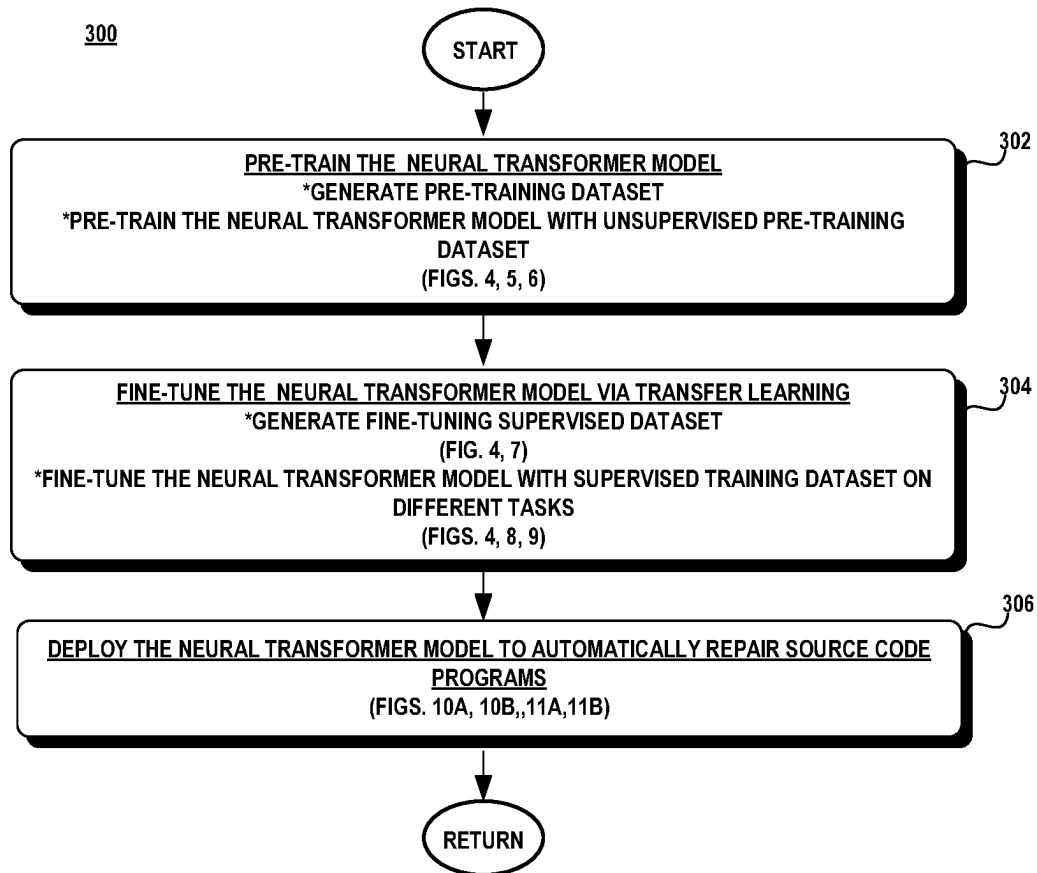


FIG. 2

**FIG. 3**

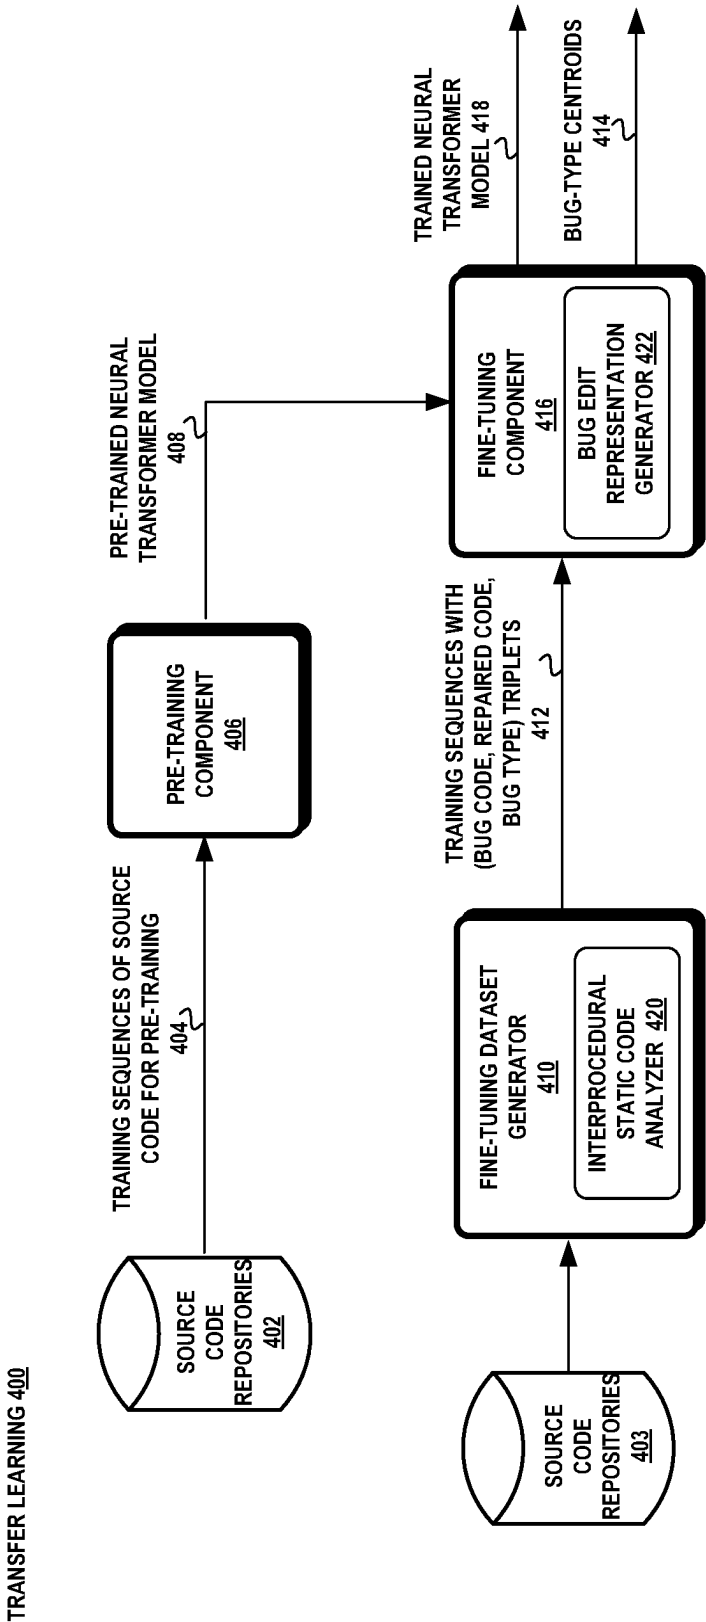
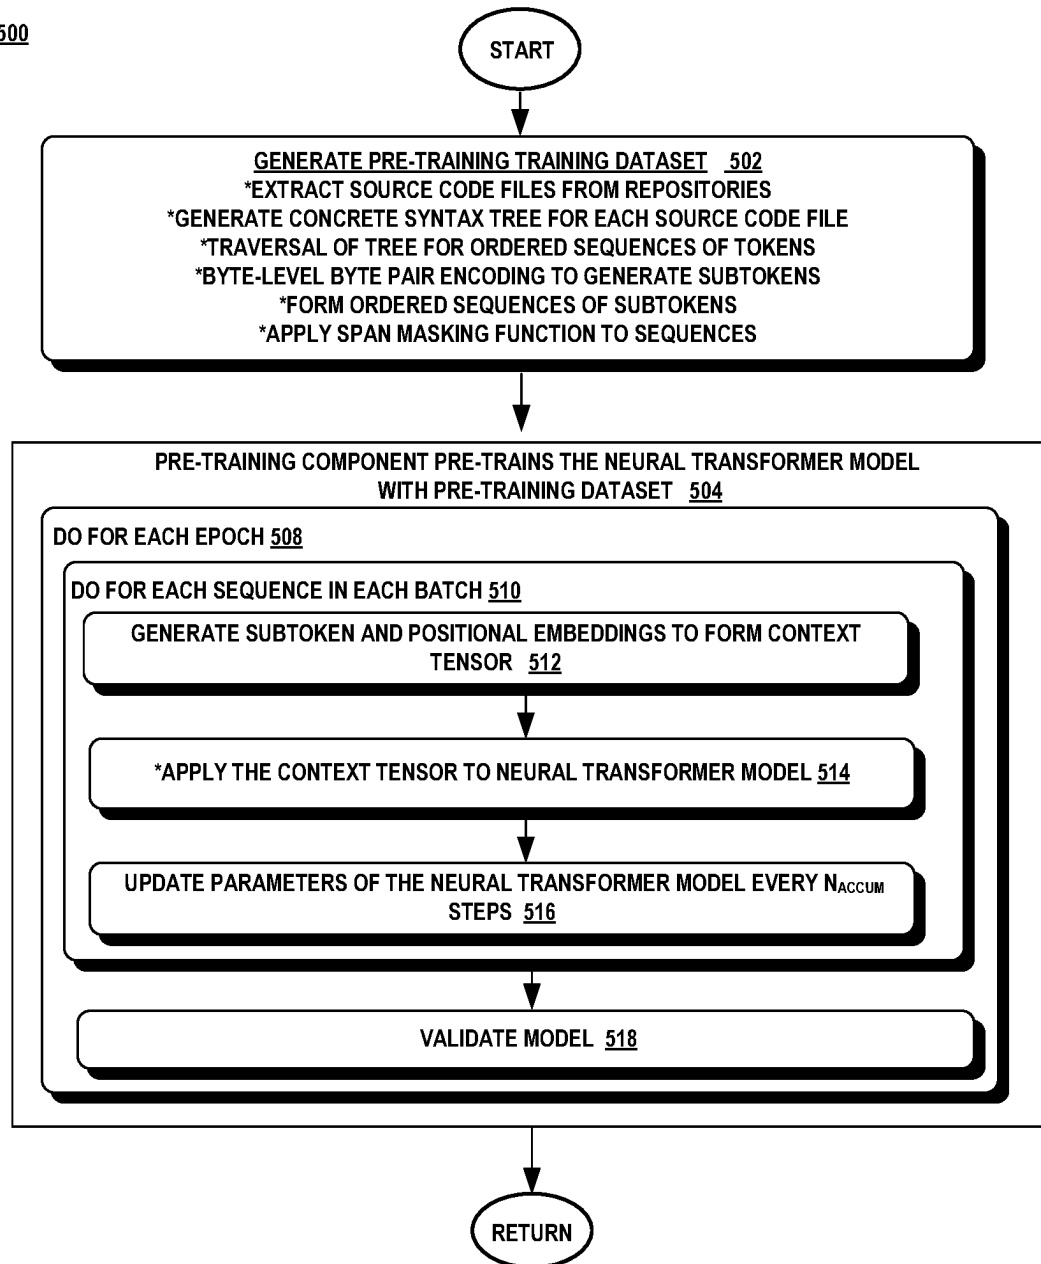


FIG. 4

500**FIG. 5**

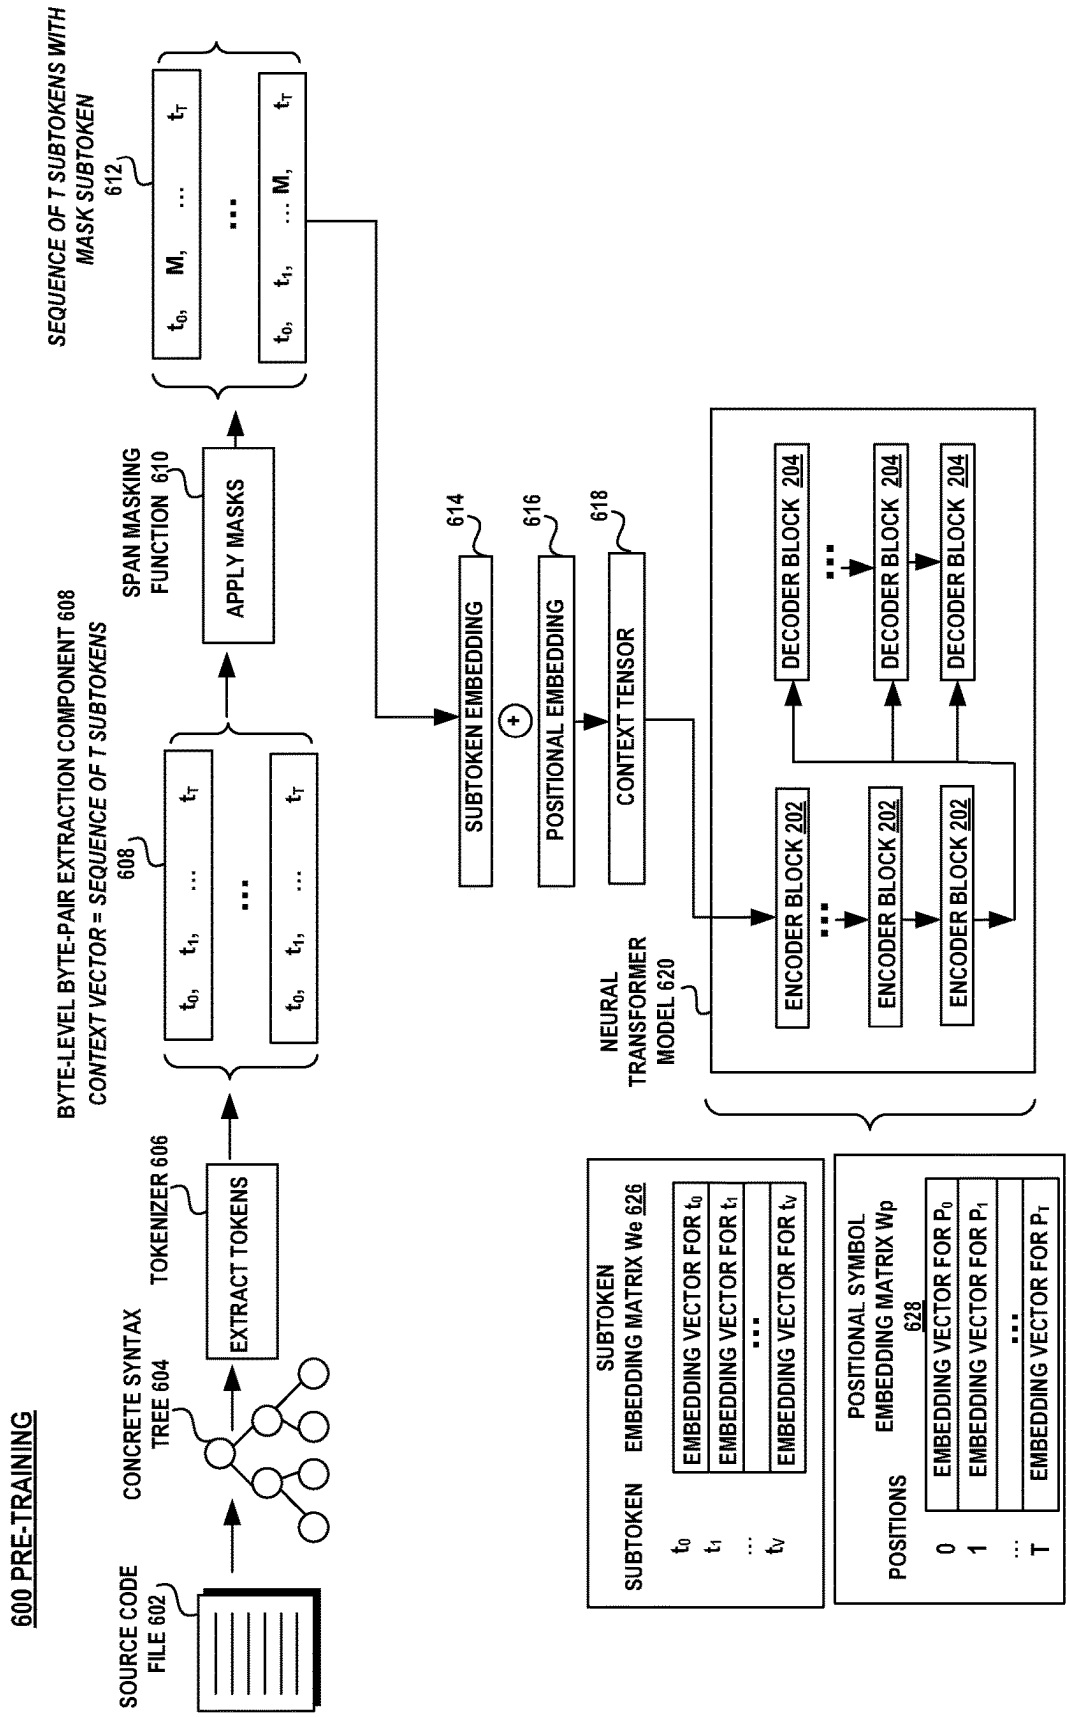
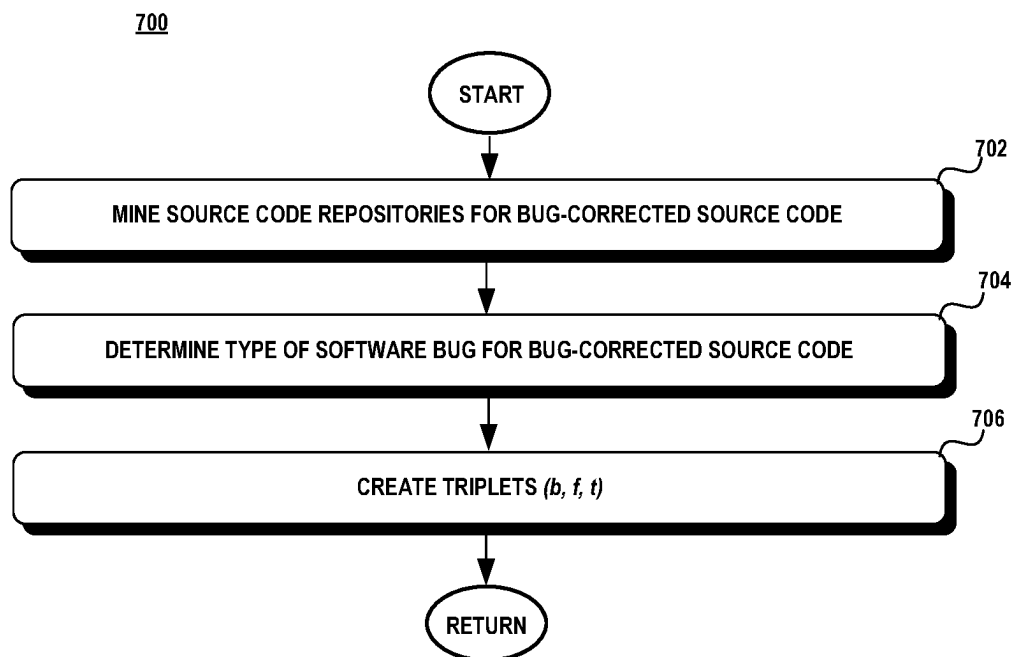
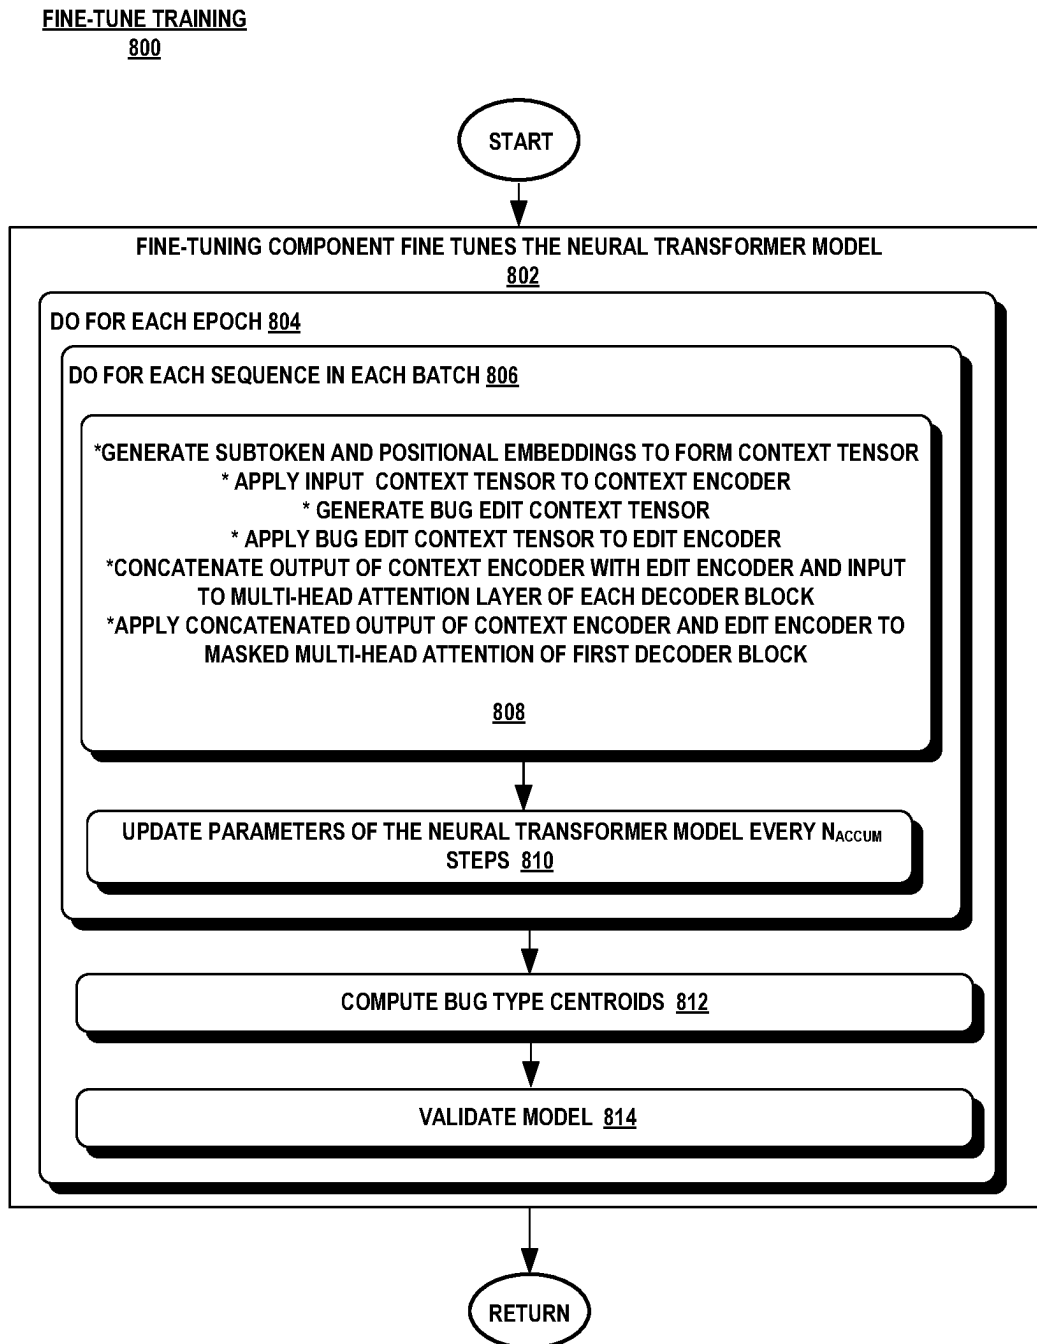


FIG. 6

**FIG. 7**

**FIG. 8**

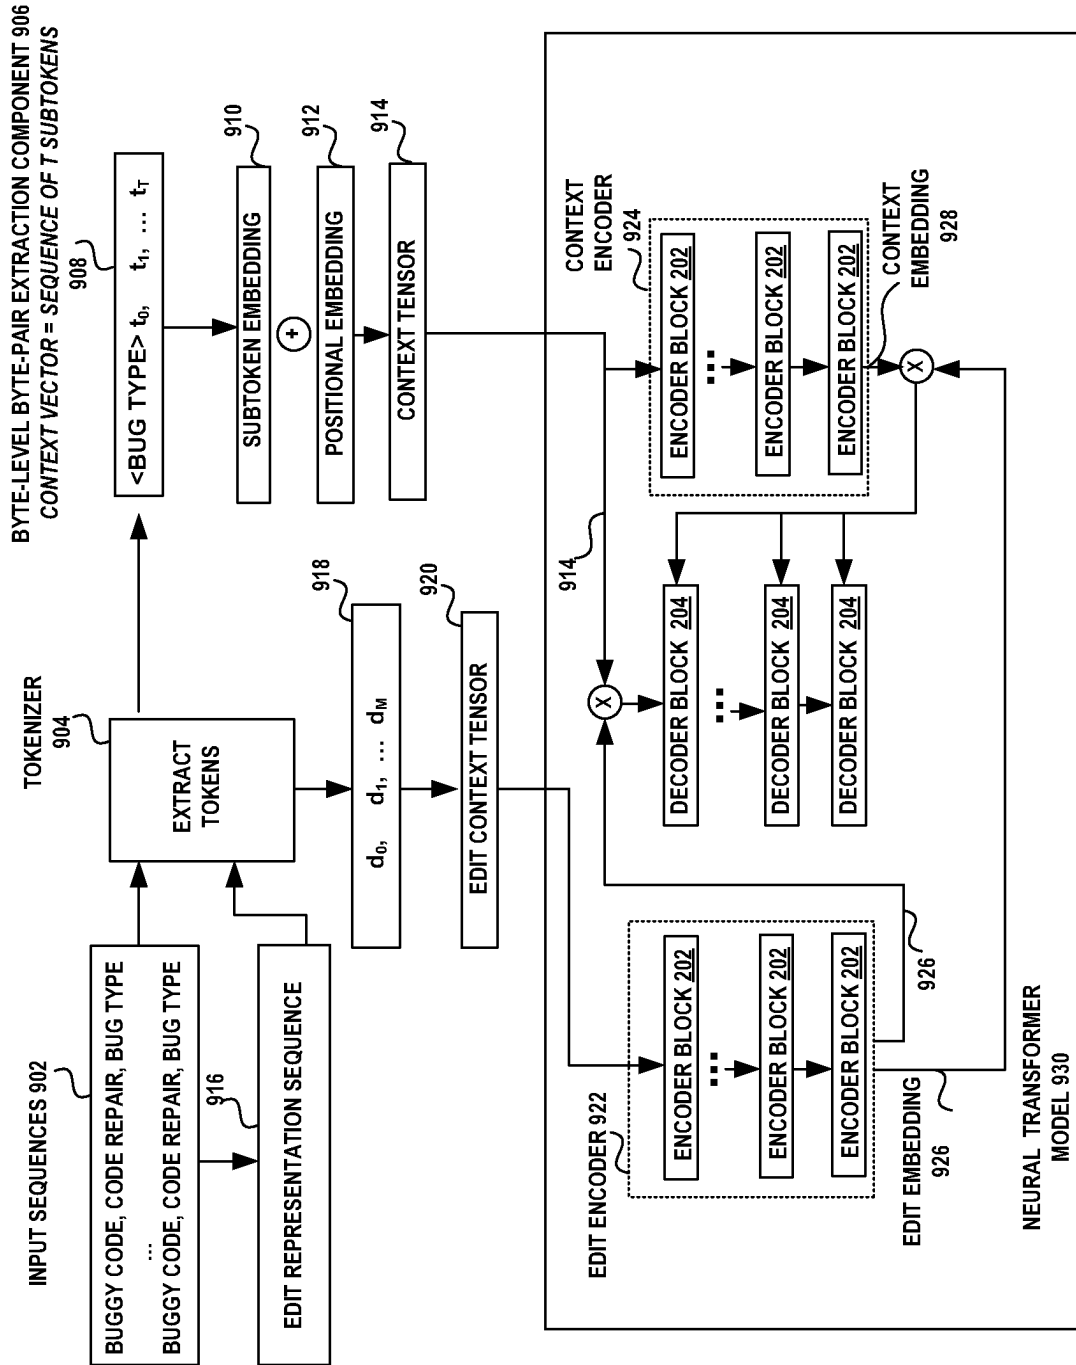
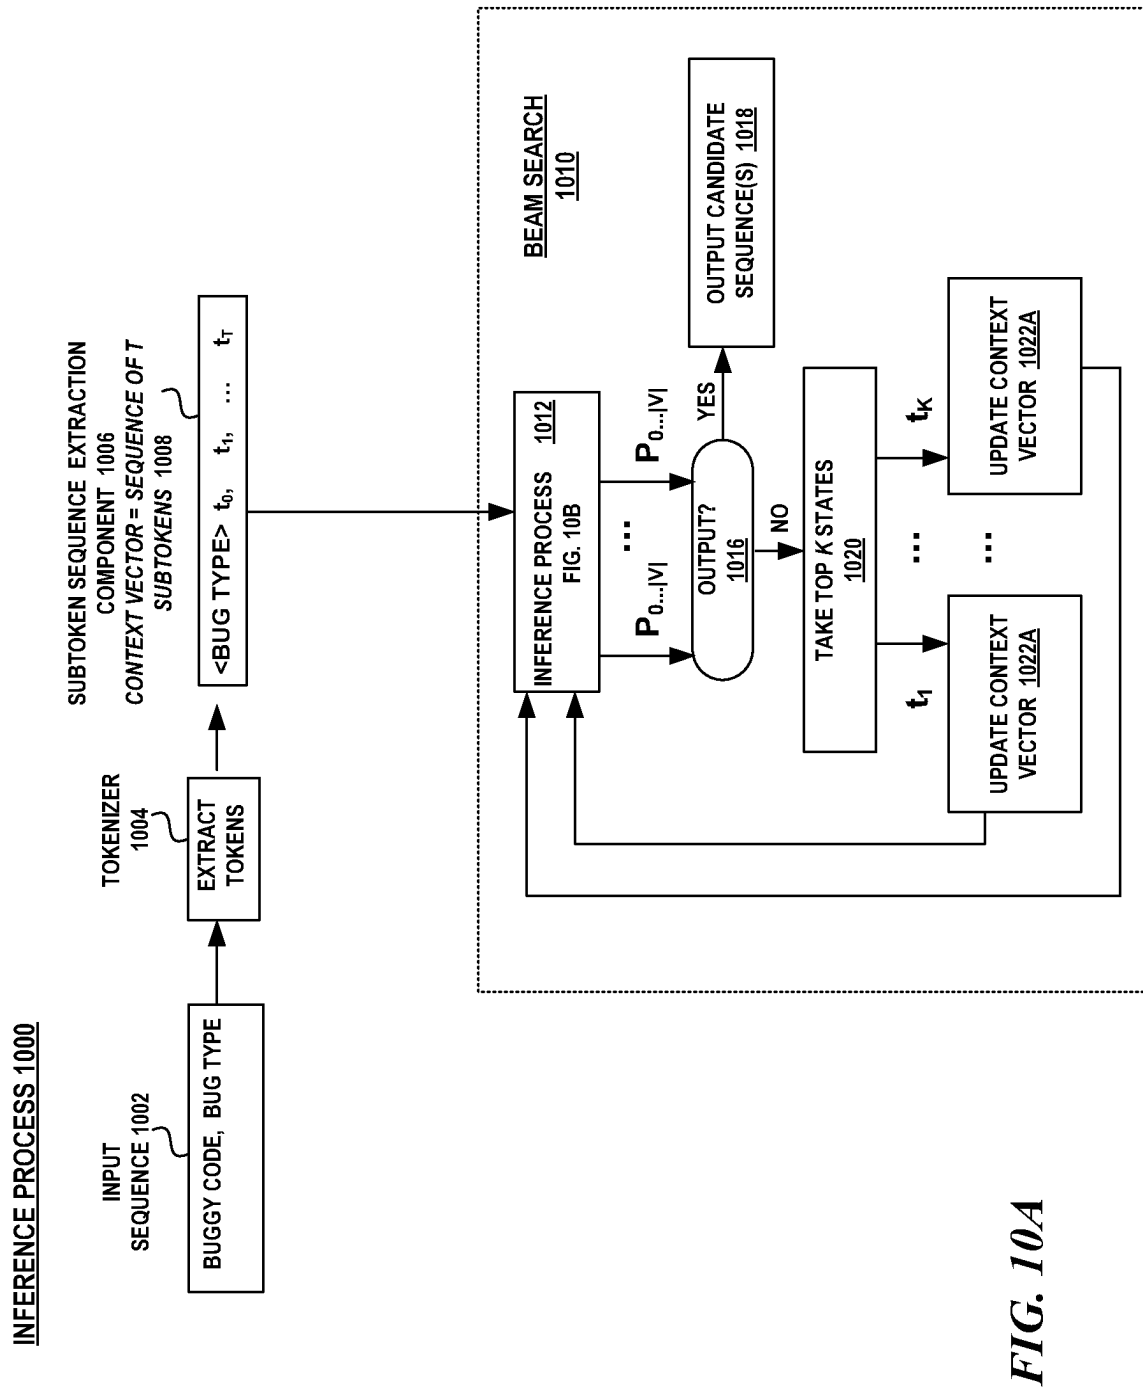
900 FINE-TUNING TRAINING

FIG. 9



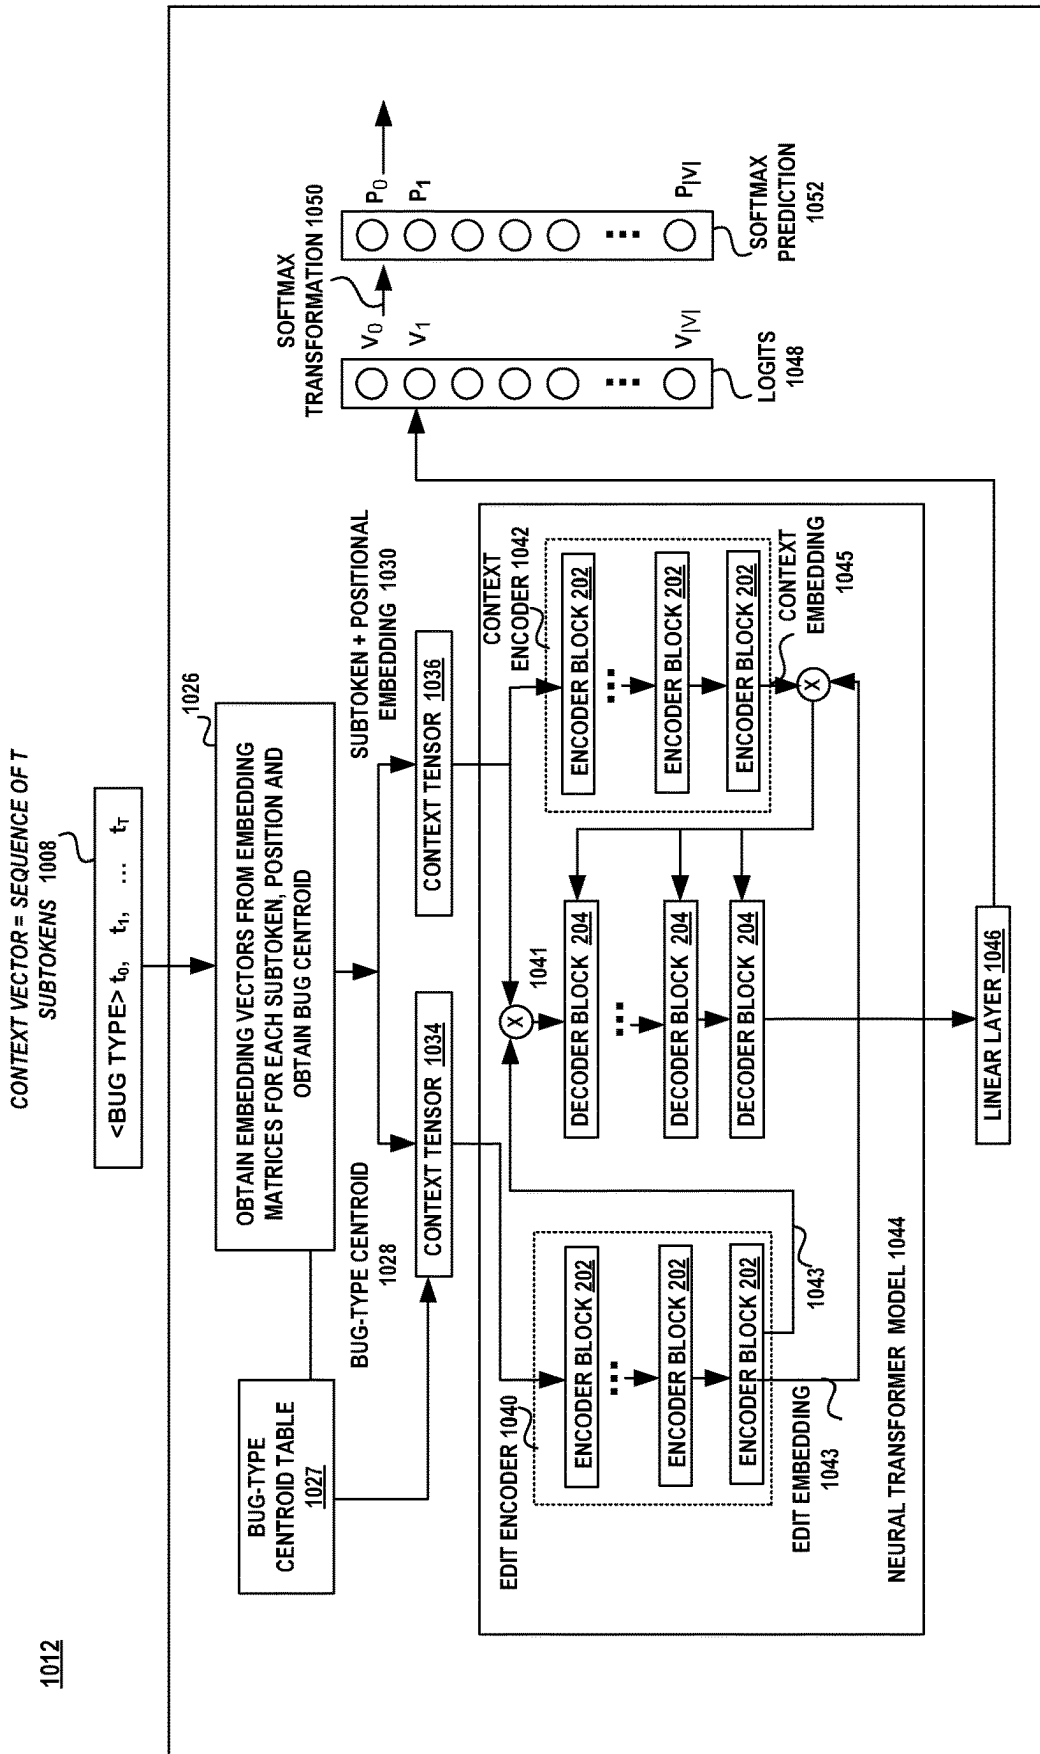
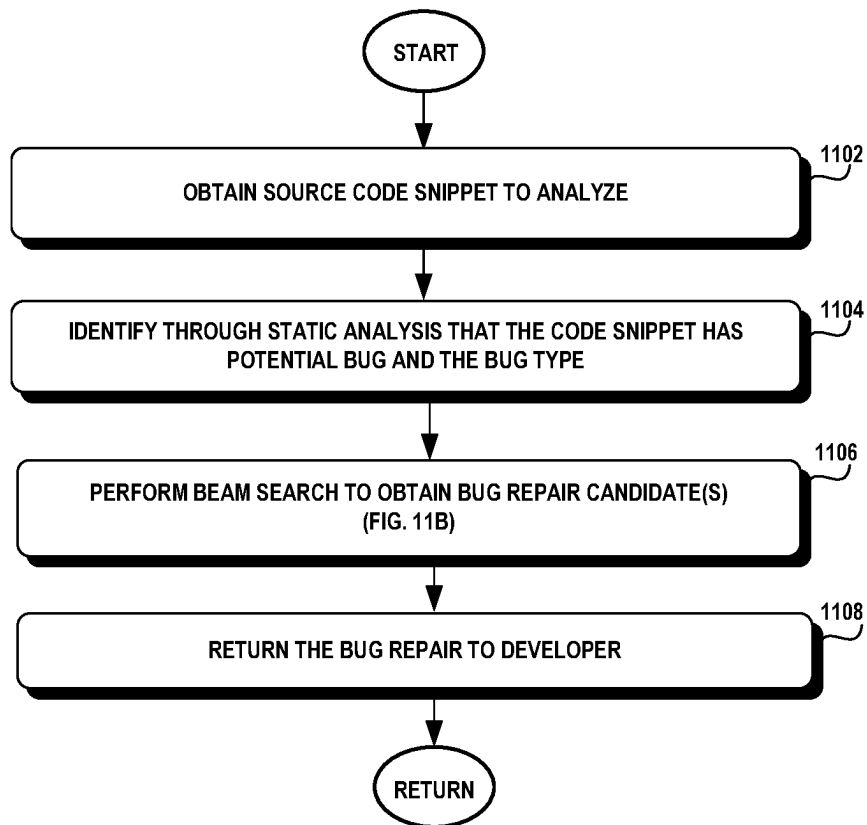
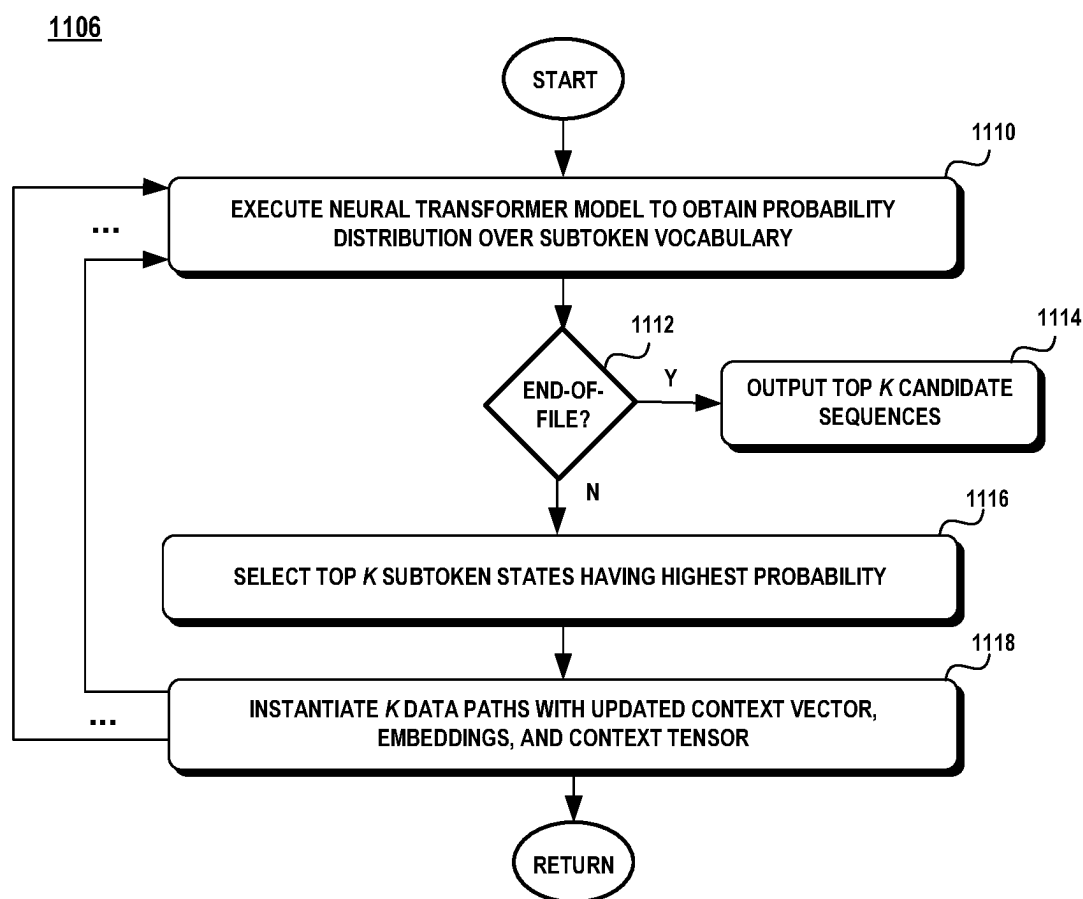


FIG. 10B

1100**FIG. 11A**

**FIG. 11B**

1200

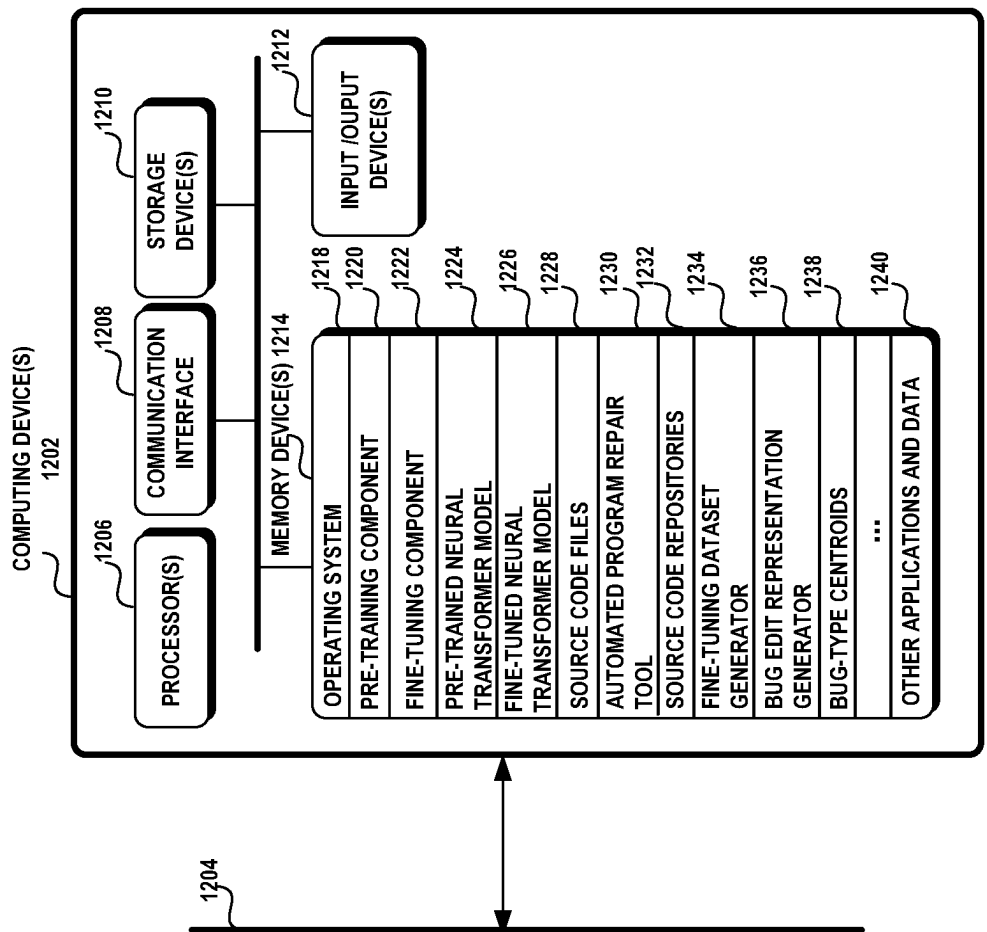


FIG. 12

AUTOMATED PROGRAM REPAIR TOOL**CROSS-REFERENCE TO RELATED APPLICATIONS**

This application is a continuation of U.S. patent application Ser. No. 17/994,185, filed on Nov. 25, 2022, which is a continuation of U.S. Pat. No. 11,526,424, filed on Jun. 10, 2020, which claims the benefit of provisional application No. 63/025,535 filed on May 15, 2020, each of which is incorporated by reference herein in its entirety.

BACKGROUND

During the development of a program or software, a range of measures is taken to ensure that the program is tested prior to the release and distribution of the program. These measures are aimed at reducing the number of bugs in the program in order to improve the quality of the program. A bug in a source code program is an unintended state in the executing program that results in undesired behavior. There are different types of software bugs which may not be detected before the program is released.

Static analysis tools are often used to detect certain types of bugs, such as syntax errors. However, static analysis tools are not adept at analyzing runtime behavior and cannot detect runtime errors. Testing is used to identify software bugs that occur at runtime. It is impossible to test all possible user scenarios and at times, the testing is limited to certain user scenarios. In addition, tests are ineffective at discovering certain unknown bugs or defects deterministically, such as resource leaks, memory leaks, null pointer dereferences, and concurrency errors, which are difficult to detect deterministically.

Software maintenance makes the corrective measures needed to fix software bugs after the bugs are reported by end users. Fixing the software bugs after deployment of the program hampers the usability of the deployed program and increases the cost of the software maintenance services.

SUMMARY

This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

An automated program repair tool is based on a sequence-to-sequence neural transformer model with attention to predict a bug repair in the context of a code snippet containing the source code bug and its identified bug type. The neural transformer model detects similar properties among certain types of source code bugs across different contexts and domains and learns specific bug-fix patterns for common bug types. Bugs belonging to the same category can be fixed using similar patterns of code changes.

The neural transformer model is pre-trained on a large unsupervised corpus of source code using a span-masking denoising optimization objective, and fine-tuned on a large supervised dataset of triplets containing a bug-type annotation, software bug, and repair. The bug-type annotation is derived from an interprocedural static code analyzer which relies on mathematical logic and symbolic reasoning to detect common bug types.

For each bug within a bug type category, an edit embedding representation is generated which aims to encapsulate

essential information of the bug type and the code changes needed to fix it. Subsequently, a single bug-type edit centroid is computed for each bug type category, from the edit embeddings of each bug of the same type. The bug-type edit centroid is then used during inference in the decoding phase to generate the bug repair for bugs belonging to the same category. Specifically, the bug type annotation and edit representation are used during fine-tuning, while the bug-type centroid is used during inference in place of the edit representation, when the bug repair is not available.

These and other features and advantages will be apparent from a reading of the following detailed description and a review of the associated drawings. It is to be understood that both the foregoing general description and the following detailed description are explanatory only and are not restrictive of aspects as claimed.

BRIEF DESCRIPTION OF DRAWINGS

FIG. 1 illustrates an exemplary automated program repair tool that predicts a code repair for a buggy source code snippet having an identified bug type.

FIG. 2 is a schematic diagram illustrating an exemplary architecture of a neural transformer model based on encoder and decoder blocks with attention.

FIG. 3 is a flow diagram illustrating an exemplary method for training the neural transformer model and deploying the neural transformer model to automatically repair buggy source code programs.

FIG. 4 is a schematic diagram illustrating an exemplary transfer learning system for training the neural transformer model.

FIG. 5 is flow diagram illustrating an exemplary method for pre-training the neural transformer model.

FIG. 6 is a schematic diagram illustrating an exemplary process for pre-training the neural transformer model and the components used therein.

FIG. 7 is a flow diagram illustrating an exemplary method for generating the fine-tuning dataset for the neural transformer model.

FIG. 8 is a flow diagram illustrating an exemplary method for fine-tuning the neural transformer model.

FIG. 9 is a schematic diagram illustrating an exemplary process for fine-tuning the neural transformer model and the components used therein.

FIGS. 10A-10B are schematic diagrams illustrating the inference process that uses the neural transformer model.

FIG. 11A-11B is a flow chart illustrating an exemplary method for generating candidate bug repair code snippets using the neural transformer model.

FIG. 12 is a block diagram illustrating an operating environment.

DETAILED DESCRIPTION**Overview**

The subject matter disclosed pertains to automated program repair based on a sequence-to-sequence neural transformer model with attention. Automated program repair is the task of predicting the contents of a software bug fix in the context of a code snippet containing a software bug and its identified bug type.

Certain types of software bugs have similar properties across different contexts and domains and can be fixed using similar patterns of code changes. The neural transformer model learns specific bug-fix patterns for common bug

types. Bugs belonging to the same category can be fixed using similar patterns of code changes.

In one aspect, the neural transformer model focuses on memory safety software bugs such as null dereference, immutable cast, empty vector access, memory leaks, and thread-safety violations. Null pointer dereference occurs when the program dereferences a pointer that it expects to be valid, but is null, or points to memory that has not been allocated. Null pointer dereferences typically cause the program to crash or exit. An immutable cast is an unsafe cast operation where it is not possible to cast a variable of one data type into another data type. For example, it is not possible to cast a null string into a non-null string.

An empty vector access error occurs when a program attempts to access a vector that has not been allocated. A race condition is a thread safety error that occurs when two threads attempt to access a shared memory address at the same time. A memory leak occurs when a program allocates memory without eventually releasing it. Eventually, the program will exhaust all the available memory and crash when the program attempts to allocate additional memory.

The neural transformer model is trained on a large unsupervised corpus of source code using a span-masking denoising optimization objective, and fine-tuned on a large supervised dataset of triplets containing a bug-type annotation, software bug, and its repaired version. The bug-type annotation is derived from an interprocedural static code analyzer which relies on mathematical logic and symbolic reasoning to detect common bug types.

For each bug within a bug-type category, an edit embedding representation is generated which aims to encapsulate essential information of the bug type and the code changes needed to fix it. Subsequently, a single bug-type edit centroid is computed for each bug-type category, from the edit embeddings of each bug of the same type. The bug-type edit centroid is then used in the decoding phase to generate the bug repair for bugs belonging to the same category. Specifically, the bug type annotation and edit representation are used during fine-tuning, while the bug-type centroid is used during inference in place of the edit representation, when the bug repair is not available.

FIG. 1 shows an exemplary automated program repair system 100 in which a program repair tool 102 receives a code snippet 104 having been identified as having a source code bug 104 and the corresponding bug type 106. The code snippet 104 in FIG. 1 is written in the Java programming language and has a line of code with an identified null pointer dereference, `if (connection.isValid(7))`. A null pointer dereference occurs when a program dereferences a pointer or value that it expects to be valid but is null. In order to avoid this problem, the program should check if the connection object is not null before invoking the `isValid` method.

As shown in FIG. 1, the program repair tool 102 provides a proposed repair 110 for the erroneous line of code in a repaired code snippet 108. The repair includes a check to ensure that the connection object is not null, `if ((connection) != null) && (connection.isValid(7))`, before attempting to invoke the `isValid()` method.

The program repair tool 102 is based on a neural transformer model with attention trained on various source code programs. The neural transformer model handles dependencies between its input and output with attention and without using recurrent neural networks (RNN) (e.g., long short-term memory (LSTM) network) and convolutional neural networks (CNN). Attention is a mechanism that identifies which parts of an input sequence are relevant to each symbol

in the output sequence and allows the neural transformer to access the entire input sequence all at once.

Attention now turns to a description of the architecture of the neural transformer model with attention.

Neural Transformer Model Architecture

FIG. 2 shows an exemplary structure of the neural transformer model in an encoder-decoder configuration. The neural transformer model 200 contains one or more encoder blocks 202 and one or more decoder blocks 204. The initial inputs to an encoder block 202 are the input embeddings 206 of an input sequence of the training dataset. In order to retain the order of the tokens in the input sequence, positional embeddings 208 are added to the input embedding 206 forming a context tensor 209. The initial inputs to the decoder block 204 are a shifted sequence of the output embeddings 218 from the previous time step to which the positional embeddings 220 are added forming context tensor 219.

An encoder block 202 consists of two layers. The first layer includes a multi-head attention component 210 followed by layer normalization component 212. The second layer includes a feed-forward neural network 214 followed by a layer normalization component 216. The context tensor 209 is input into the multi-head attention layer 210 of the encoder block 202 with a residual connection to layer normalization 212. The output of the layer normalization 212 is input to the feed forward neural network 214 with another residual connection to layer normalization 216. The output of the encoder block 202 is a set of hidden representations 217. The set of hidden representations 217 is then sent through additional encoder blocks, if multiple encoder blocks exist, or to the decoder 204.

Attention is used to decide which parts of the input sequence are important for each subtoken, especially when decoding long sequences since the encoder is limited to encoding a fixed-size vector. Attention mechanisms gather information about the relevant context of a given subtoken and then encode that context into a vector which represents the subtoken. It is used to identify the relationships between subtokens in the long sequence while ignoring other subtokens that do not have much bearing on a given prediction.

The multi-head attention component 210 takes a context tensor 209 and weighs the relevance of each subtoken represented in the context tensor to each other by generating attention weights for each subtoken in the input embedding 206. In one aspect, the attention function is scaled dot-product attention which is described mathematically as follows:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V,$$

where the input consists of queries Q and keys K of dimension d_k , and values V of dimension d_v . Q is a matrix that contains the query or vector representation of one subtoken in a sequence, K is the vector representations of all subtokens in the sequence, and V is the vector representations of all the subtokens in the sequence.

The queries, keys and values are linearly projected h times in parallel with d_v output values which are concatenated to a final value:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O,$$

$$\text{where } \text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V),$$

5

with parameter matrices $W_i^Q \in \mathbb{R}^{d_{model} \times d_k}$, $W_i^K \in \mathbb{R}^{d_{model} \times d_k}$, $W_i^V \in \mathbb{R}^{d_{model} \times d_k}$, and $W^O \in \mathbb{R}^{h_{d_o} \times d_{model}}$.

In order to reduce the training time of the neural transformer, layer normalization is used between the layers. The layer normalization component normalizes the inputs across the features. The mean and standard deviation is computed across the feature dimensions. There is a first layer normalization **212** that precedes the feed forward neural network **214** and a second layer normalization **216** that follows the feed forward neural network **214**.

The feed-forward neural network **214** processes each output encoding separately **213**. The output of the top encoder block is a set of attention vectors K and V **217** which is used by the encoder-decoder multi-head attention layer **226** of the decoder block **204**.

The decoder block **204** predicts each subtoken t_i in the target language one-by-one at each time step conditioned on all previously-generated target subtokens t_1, \dots, t_{i-1} . The decoder block **204** consists of three layers. The first layer includes a masked multi-head attention component **222** followed by a layer normalization component **224**. The output of the layer normalization component **224** is input into the encoder-decoder multi-head attention component **226** with a residual connection to layer normalization component **228**. The second layer includes an encoder-decoder multi-head attention component **226** followed by a layer normalization component **228**. The output of layer normalization component **228** is input into the feed forward neural network **230** with a residual connection to layer normalization component **232**. The third layer includes a feed forward neural network **230** followed by a layer normalization component **232**.

The masked multi-head attention component **222** receives the output embeddings of the previous timestep. The masked multi-head attention component **222** masks the output embeddings from future time steps. The encoder-decoder multi-head attention layer **226** receives queries from the previous decoder layer **225** and the memory keys and values **217** from the output of the encoder block **202**. In this manner, the decoder block **204** can attend to every position of the input sequence. The feed-forward neural network **230** processes each output encoding separately. A layer normalization component **224**, **228**, **232** is used between the layers in order to normalize the inputs across the features.

The linear layer **234** projects the vector produced by the stack of decoders into a logits vector. The softmax layer **236** then turns the scores of the logits vector into probabilities for each subtoken in the vocabulary which are positive and normalized **238**.

In one aspect, the neural transformer model contains a stack of six encoder blocks and a stack of six decoder blocks which are aggregated into a neural transformer block. The output of each encoder block is passed onto the next encoder block and processed. Each decoder block receives the attention weights computed from the last encoder block. The use of multiple stacked encoder blocks and decoder blocks increases the model's capacity allowing the model to learn increasing levels of abstraction.

Neural Transformer-Based Program Repair Generation

FIG. 3 is a flow diagram illustrating an exemplary process of a neural transformer model-based automated program repair tool **300**. Initially, the neural transformer model is trained through a transfer learning process that includes pre-training the neural transformer model with an unsupervised training dataset of source code (block **302**) and fine-tuning the neural transformer model with a supervised training dataset of translation tasks (block **304**).

6

The unsupervised training dataset includes source code snippets for the neural transformer model to learn statistical properties of the source code, such as syntactic rules of the programming languages, as well as semantic information from co-occurrence of specific variable and method names. The pre-trained model represents a base which is subsequently fine-tuned on bug repair translation tasks. The supervised training data includes triplets consisting of a buggy source code snippet, its repair code snippet, and its bug type which train the neural transformer model to learn to translate buggy code of a particular bug type into a specific bug repair. When the model has been trained and verified successfully, the model is deployed in an automatic program repair tool (block **306**).

Transfer Learning

The neural transformer model is trained through transfer learning. Transfer learning is a methodology of training models by pre-training the model using unsupervised learning on unlabeled data to learn generalized knowledge and then fine-tuning the model via supervised learning on labeled data. The neural transformer model is pre-trained on a large unsupervised training dataset of unlabeled source code that contains lines of source code in various programming languages (e.g., Python, C#, JavaScript and TypeScript) using a denoising objective and then separately fine-tuned on translation tasks.

FIG. 4 illustrates the transfer learning system **400** to train a neural transformer model with attention **418**. Turning to FIG. 4, a pre-training component **406** generates an unsupervised training dataset **404** from source code files from various source code repositories **402**. The pre-training component **406** trains the pre-trained neural transformer model **408** which is then fine-tuned by the fine-tuning component **416**. The fine-tuning dataset generator **410** generates a training dataset of triplets **412** that includes a code snippet with a bug, the repaired code snippet and a type of the bug. The fine-tuning dataset generator **410** obtains the buggy code snippets from a source code repository having repaired source code **403**.

The fine-tuning dataset generator **410** uses an interprocedural static code analyzer **420** to classify a bug type. The fine-tuning component **416** includes a bug edit representation generator **422** to compute an edit embedding representation for the bug during training, which will be replaced with a bug centroid for each bug type during inference, when the bug repair is not available.

A bug fix or repair is represented by the triplet $bf = \{b, f, t\}$, where b is the buggy code, f is the bug repair, and t is the type of bug that was fixed. Source code with a bug is obtained from a version-controlled source code repository **403**. The fine-tuning dataset generator **406** analyzes the source code repository **403** for changes made to a repository in order to identify the bugs introduced or fixed in a commit. A commit adds the latest changes made to a source code file to the repository. The files involved in the changed code are identified and input into a static analyzer to identify the bug type t . The bug type, the buggy code and the repaired code are extracted to form the triplet $bf = \{b, f, t\}$.

The fine-tuning component **416** trains the pre-trained neural transformer model **408** with a large supervised training dataset of triplets **412**. The triplets (b, f, t) represent translation tasks that teach the model to learn to translate an input sequence of buggy code and its bug type into an output sequence that contains the repaired code.

The fine-tuning component **416** also generates a bug-type edit representation for each bug type. A bug-type edit representation is a vector representation of the edits per-

formed to generate a bug fix for a certain bug type. A developer performs a sequence of edits to transform the code **b** into the code **f**. The bug-type edit representation is a vector representation of the edits that transforms the code **b** into the code **f**. A representation function **G** maps an edit operation $b \rightarrow f$ to an embedding vector $G(b, f) \in \mathbb{R}^d$, where d is the embedding dimension.

Given an edit representation function **G** and a triplet, (**b**, **f**, **t**), clusters are identified in the embedding space for each bug type. For each bug-type cluster, a bug-type embedding is generated as a centroid vector $g(\text{type})$ **414**. The centroid embedding for a particular bug type is used to inform the neural transformer model during the inference process, when predicting a bug repair.

Pre-Training

FIGS. **5** and **6** illustrate an exemplary method for pre-training the neural transformer model. Turning to FIGS. **4**, **5** and **6**, the pre-training training component **406** generates a training dataset to pre-train the neural transformer model (block **502**). The pre-training component **406** generates a pre-training dataset from a diverse corpus of unlabeled source code programs or files **602**. This is referred to as unsupervised learning since the model draws inferences from the input data without labeled responses. The pre-training component **406** extracts selected source code files **602** from various source code repositories **402**. The source code files **602** contain context beyond method bodies, method signatures, and docstrings, such as imports, globals, comments, and scripts. (Collectively, block **502**).

A source code repository **402** may be a file archive and web hosting facility that stores large amounts of source code either privately or publicly. A source code repository **402** can be structured as a version control system, such as GIT, Mercurial, etc. The source code files residing in the source code repository **402** vary and may be written in different programming languages. The selected source code files **602** can come from different domains, such as without limitation, scientific computing, web development, dataflow programming, machine learning, and the like. (Collectively, block **502**).

The pre-training component **406** transforms each of the selected source code files **602** into a concrete syntax tree **604**. The concrete syntax tree **604** represents the source code text in the parsed form. The concrete syntax tree **604** may also be a parse tree. A concrete syntax tree **604** represents the syntactic structure of a program in a hierarchical or tree structure. The concrete syntax tree **604** is an n -ary tree data structure that includes nodes that represent a construct in the grammar of the programming language of a program. The concrete syntax tree **604** includes one root node, multiple internal nodes, and multiple terminal nodes. The terminal nodes represent the tokens. A token is a symbol that represents an operand or an operator. The concrete syntax tree **604** differs from an abstract syntax tree where the terminal nodes represent operands. (Collectively, block **502**).

The pre-training component **406** uses a tokenizer **606** to extract tokens from the concrete syntax tree **604**. The frequently-used elements in a programming language are encoded into tokens and the less frequently-occurring elements are encoded into combinations of characters referred to as subtokens. For simplicity, the term subtoken shall include tokens and subtokens. (Collectively, block **502**).

The pre-training component **406** uses a byte-level byte-pair extraction algorithm **608** to generate T-ordered sequences of subtokens **609**, where T is the maximum context length. Byte-level byte-pair encoding (BBPE) is used to generate the vocabulary used by the neural transformer

model. A text string, either a sequence of source code or a natural language text, is represented as a sequence of Unicode Transform Format, UTF-8 bytes. The input text string of subtokens is encoded as a sequence of UTF-8 bytes, where a subtoken is encoded into one to four bytes. A byte sequence is then partitioned into byte-level subwords, referred to as byte n -grams. (Collectively, block **502**).

The byte-level subwords are generated using the Byte Pair Encoding (BPE) algorithm, which extracts the k most frequently-occurring n -grams. The result is a vocabulary size of the k most frequently-occurring n -grams. An n -gram is a contiguous sequence of n subtokens from an input text string of either source code or natural language text. This type of encoding does not rely on knowing the underlying language making it suitable for an input sequence of text strings that contain source code or natural language text. The ordered sequences of UTF-8 bytes are translated into a T-ordered sequence of subtokens which are vector representations of a source code fragment or natural language text. The T-ordered sequence of subtokens are represented in a context vector **610**. (Collectively, block **502**).

A denoising function, such as a span masking function **610**, is then applied to each sequence that randomly masks out a subset of subtokens and the masked span of subtokens is replaced with a mask subtoken, **M**. The model is trained with the masked sequences to learn to reconstruct the original sequence without the masked subtokens. In one aspect, the mask subtoken replaces a span of subtokens. The number of text spans and the span lengths are randomly generated and each span is replaced with a single mask subtoken. The masked denoising is based on the cloze task of evaluating human language-learners' proficiency, in which humans are given a foreign language with missing words, and are asked to correctly choose the missing word. The benefit of span-masking denoising in pre-training is that the model learns the desired language in an unsupervised fashion, but also is bi-directional in the sense that it learns the relationships of words both before and after their occurrence. (Collectively, block **502**).

The pre-training component **406** then pre-trains the neural transformer model with the pre-training dataset **404** (block **504**). Neural transformer models are trained iteratively, making multiple passes over the training dataset before converging to a minimum. An epoch represents the entire training dataset passed forwards and backwards through the neural transformer block once. Since the training dataset is very large, it is partitioned into smaller batches. The training is iterative and the entire dataset is passed through the neural transformer in multiple iterations. Each training iteration includes forward propagation, loss calculation, backpropagation steps followed by updating the weights. The training dataset is partitioned into batches with each batch of sequences running through the training process. (Collectively, block **504**).

The neural transformer model has multiple blocks and layers so that more detailed relationships within the data are learned as well as how the features interact with each other on a non-linear level. The model architecture, training procedure, data normalization and vocabulary encoding procedures are hyperparameters that are tailored to meet a particular objective. The values of the hyperparameters influence how the parameters are learned.

In one aspect, the hyperparameters may include the following: (1) subtoken and position embedding layers of dimensions: 30000×768 , and 1024×768 respectively; (2) the configuration of the neural transformer model with six encoder blocks and six decoder blocks; (3) for the training

procedure: denoising auto-encoder, with a cross-entropy loss optimization objective; the sequence length of 1024 symbols; a mini-batch size of 8; the gradient accumulation steps for each weight update is 8; the Adam stochastic optimization procedure is used to train the feed forward neural network; and the learning rate is 0.0001; (4) the data normalization procedure: normalize all string and numerical literals, keeping the ten most frequent; and (5) the vocabulary encoding procedure: byte-level byte-pair encoding, preserve the ten most frequent string and numerical literals encoding them as a single token during byte-level byte-pair encoding procedure; and introduce special control flow tokens to denote end-of-line, end-of-file, end-of-method, end-of docstring, dedent, and indent symbols.

For each sequence of each batch in each epoch (blocks 508, 510), the T-ordered sequences of subtokens are then mapped into numeric vectors and then into respective subtoken embeddings 614 and positional embeddings 616 (block 512). An embedding is a learned representation for the text-based subtokens where subtokens that have a common meaning have a common representation. An embedding is a mapping of discrete categorical variables to a vector of continuous numbers. There is an embedding for each subtoken in the vocabulary 614 and a corresponding positional embedding 616. The subtoken embedding 614 represents the learned representation for the subtoken. The neural transformer model does not read each subtoken sequentially and as such, has no knowledge of the subtoken's position in a sequence without additional position information. The positional embedding 616 is used to embed position information about a subtoken's position in a sequence into the neural transformer model 620.

Initial values are generated for the subtoken embedding and positional embeddings of each sequence which are then used to form a context tensor 618. Thereafter, the neural transformer model 620 learns the values for each embedding. Upon the completion of the training phase, the embeddings for each subtoken and the positional embeddings are saved into respective matrices for later use. There is a subtoken embedding matrix, We , 626 that contains an embedding vector for each subtoken t_i , $i=0 \dots V$, and a positional embedding matrix, Wp , 628 that contains an embedding vector P_j , $j=0 \dots T$, for each position, where V is the size of the vocabulary and T is the length of the subtoken sequence. (Collectively, block 512).

The encoder blocks 202 of the neural transformer model 620 take the context tensor 618 as input and passes it through the multiple blocks of multi-head attention and feed-forward neural network to finally produce a contextualized token representation. The decoder blocks 204 of the neural transformer model 620 takes a shifted sequence of an output embedding as input. (Collectively, block 514).

The feed forward neural networks in the encoder blocks 202 and the decoder blocks 204 are trained iteratively, making multiple passes over the training dataset before converging to a minimum. Each training iteration includes forward propagation, loss calculation, backpropagation steps followed by updating the weights by calculating the weight gradients. The loss function estimates the loss or error which is used to compare how good or bad the predicted results are. In one aspect, a categorical cross-entropy loss function is used. Once the loss is calculated, it is propagated backwards to the hidden layer that contributed directly to the output. In backpropagation, the partial derivatives of the loss function with respect to the trainable parameters are determined. The weight gradients are calculated as the difference between the old values and the new

values of the weights. The weights are adjusted to make the loss as small as possible using a gradient descent technique. In one aspect, a Stochastic Gradient Descent (SGD) method is the optimization algorithm used to find the values of parameters of the function that minimizes the loss function. A backpropagation through time (BPTT) algorithm may be used to update the weights. (Collectively, block 514).

At the completion of each batch, the parameters of the neural transformer model are updated at a preconfigured frequency denoted as Naccum. Naccum is a gradient accumulation frequency and in one aspect has a value of 8. The parameters include the subtoken embeddings and the positional embeddings which are stored in a respective embedding matrix. (Collectively, block 516).

Next, the neural transformer model is validated. Before the neural transformer model is trained, a set of hyperparameters is selected randomly and then tuned to achieve a desired performance. The neural transformer model is tested using a validation dataset to determine the appropriate hyperparameters settings to achieve a desired goal. When the desired goal is not achieved, one or more hyperparameters are adjusted and the training is repeated until the target goal is achieved. Perplexity on the validation set is calculated to validate the performance of the model with respect to the learning the masked out original text. (Collectively, block 518).

Fine-Tuning

After the pre-training is performed, the neural transformer model is fine-tuned on individual supervised tasks. The fine-tuning training dataset is generated from software bugs and repairs from various programs in different programming languages. A static analyzer is used to identify the type of software bug. For each bug type category, a canonical edit embedding representation is generated which aims to encapsulate essential information on the bug type and the code changes needed to fix it. A bug type edit centroid is computed for each bug type from the canonical edit embedding of each bug type and used during inference in the decoding phase to generate the bug repair.

The fine-tuning ensures that the neural transformer model understands the relationship between a code snippet with bugs, the bug type, and its repair by training the model with sequences of these features. The repetition of the features improves the model's ability to learn the relationships between the different features. Bug-fixing edits carry meaningful semantic information about bugs, and are thus capable of improving accuracy of the program repair system. During the fine-tuning training stage, the exact bug-fixing edit for a given buggy code snippet and a fixing code snippet is used. For a given training sample, the edit embedding vector is concatenated with the encoder output vector and is utilized to initialize the decoder state, it is also concatenated with output embedding at each temporal step. However, the exact bug-fixing edits are not known ahead of time during inference. Given that embedding vectors of bug-fixing edits belonging to a given bug type are geometrically close and are forming clusters in the embedding space, centroids of these clusters are used as an alternative way of passing semantic information about bug-fixing edits to the model decoder during inference time.

Turning to FIGS. 4 and 7, the fine-tuning dataset generator 410 mines various source code repositories 403 for source code that has been corrected due to a software bug. The source code file containing the original code with the source code bug and the source code file containing the bug repair are obtained from the source code repository. (Collectively, block 702).

The fine-tuning dataset generator **410** uses a static code analyzer to determine the type of software bug for the source code file having the bug. In one aspect, an interprocedural static analyzer **420** is used to analyze a source code file to determine whether the source code has a software bug and to identity the type. An interprocedural static code analyzer **420** analyzes program code without executing the program. The analysis is interprocedural rather than intraprocedural. An intraprocedural analysis is performed within a method, otherwise referred to as a procedure or function. An interprocedural analysis spans multiple files or methods, including all the methods in the entire program. An interprocedural static code analysis is able to detect memory safety faults or bugs that span multiple files or methods, such as null pointer dereferencing and memory leaks, which would be missed if intraprocedural static analysis were used. (Collectively, block **704**).

In one aspect, the interprocedural static code analyzer **420** is Facebook®'s Infer. Infer identifies memory safety errors, such as, null pointer dereference, memory leak, immutable cast, empty vector access, and thread safety violation. Infer is based on separation logic proof rules in order to discover program paths with the symbolic heap that violate heap-based properties. In this manner, Infer is able to detect a memory safety-type program error and to identify the particular type. (Collectively, block **704**).

The fine-tuning dataset generator **410** forms training sequences of triplets including the code snippet containing the software bug, the code snippet of the repair, and the bug type identified from the static analyzer (block **706**).

Fine-Tuning Training

Referring to FIGS. **4**, **8** and **9**, the fine-tuning component **416** trains the neural transformer model and an edit representation function G with input sequences **902** having a code snippet with buggy code (buggy code), the associated code repair (code repair), and the associated bug type (bug type). The fine-tuning dataset consists of bugs and fixes, and associated bug types. A bug-fixing edit representation is a vector representation of the edits performed to generate a bug fix for a certain bug-fix pair. A developer performs a sequence of edits to transform the code b into the code f . The bug-fixing edit representation is a vector representation of the edits that transforms the code b into the code f . A representation function G maps an edit operation $b \rightarrow f$ to an embedding vector $G(b, f) \in \mathbb{R}^d$, where d is the embedding dimension. (Collectively, block **802**).

Given an edit representation function G and a triplet, (b, f, t) , clusters are identified in the embedding space for each bug type. For each bug-type cluster, a bug-type embedding is generated as a centroid vector $g(\text{type})$ **414**. The centroid embedding for a particular bug type is used to inform the neural transformer model during the inference process. The edit representation function G is the edit encoder **922** of fine-tuned model **416**, **420**. The bug edit representation generator **422** uses the edit encoder **922**. (Collectively, block **802**).

A tokenizer **904** generates a concrete syntax tree for the buggy code and the repaired code from which an ordered sequence of tokens is extracted **902**. The tokens are then converted into a byte-level byte pair encoded representation using a byte-level byte-pair extraction component **906** and

then into an ordered sequence of subtokens **908**. The input sequence **908** is prepended with the bug type, $\langle \text{BUG TYPE} \rangle$. (Collectively, block **802**).

Neural transformer models are trained iteratively, making multiple passes over the training dataset before converging to a minimum. An epoch represents the entire training dataset passed forwards and backwards through the neural transformer block once. Since the training dataset is very large, it is partitioned into smaller batches. The training is iterative and the entire dataset is passed through the neural transformer in multiple iterations. Each training iteration includes forward propagation, loss calculation, backpropagation steps followed by updating the weights. The training dataset is partitioned into batches with each batch of sequences running through the training process.

The neural transformer model has multiple blocks and layers so that more detailed relationships within the data are learned as well as how the features interact with each other on a non-linear level. The model architecture, training procedure, data normalization and vocabulary encoding procedures are hyperparameters that are tailored to meet a particular objective. The values of the hyperparameters influence how the parameters are fine tuned for the particular translation task.

For each sequence of each batch in each epoch (blocks **804**, **806**), the T-ordered sequences of subtokens are then mapped into numeric vectors and then into respective subtoken embeddings **910** and positional embeddings **912** to form a context tensor (block **808**). Initial values for the subtoken embedding and positional embeddings of each sequence were learned in the pre-training and are now updated based on the input sequences (block **808**).

An edit representation sequence **916** is generated to include the edits made to the buggy source code to produce the bug repair. The tokenizer **904** extracts tokens from the edit representation sequence **916**, to form a tokenized sequence, d_0, d_1, \dots, d_M , **918** which forms the edit context tensor **920**. The edit context tensor **920** is input into an edit encoder **922** which learns an associated edit embedding **926**. The encoder blocks of the edit encoder take the edit context tensor as input and passes it through multiple blocks of multi-head attention and feed-forward neural network in each encoder layer to finally produce an edit embedding **926**. (Collectively, block **808**).

The encoder blocks **202** of the neural transformer model **920** take the context tensor **914** as input and pass it through the multiple blocks of multi-head attention and feed-forward neural network in each encoder layer to finally produce a context embedding **928**. The context embedding **928** is concatenated with the edit embedding **926** and input into each multi-head attention layer of each decoder block. The initial decoder block **204** of the neural transformer model **920** takes a shifted sequence of an output embedding from the previous time step which is concatenated with the edit embedding **926**. (Collectively, block **808**).

The feed forward neural networks in the encoder blocks **202** and the decoder blocks **204** are trained iteratively, making multiple passes over the training dataset before converging to a minimum. Each training iteration includes forward propagation, loss calculation, backpropagation steps followed by updating the weights by calculating the weight gradients. The loss function estimates the loss or error which is used to compare how good or bad the predicted results are. In one aspect, a categorical cross-entropy loss function is used. Once the loss is calculated, it is propagated backwards to the hidden layer that contributed directly to the output. In backpropagation, the partial deriva-

13

tives of the loss function with respect to the trainable parameters are determined. The weight gradients are calculated as the difference between the old values and the new values of the weights. The weights are adjusted to make the loss as small as possible using a gradient descent technique. In one aspect, a Stochastic Gradient Descent (SGD) method is the optimization algorithm used to find the values of parameters of the function that minimizes the loss function. A backpropagation through time (BPTT) algorithm may be used to update the weights. (Collectively, block **808**).

At the completion of each batch, the parameters of the neural transformer model are updated at a preconfigured frequency denoted as Naccum. Naccum is a gradient accumulation frequency and in one aspect has a value of 8. The parameters include the subtoken embeddings and the positional embeddings which are stored in a respective embedding matrix. (Collectively, block **810**).

In addition, the bug type centroids are computed for each bug type and stored in a bug centroid table for use during the inference phase (block **812**).

Next, the neural transformer model is validated. The neural transformer model is tested using a validation dataset to determine the appropriate hyperparameters settings to achieve a desired goal. When the desired goal is not achieved, one or more hyperparameters are adjusted and the training is repeated until the target goal is achieved. Perplexity on the validation set is calculated to validate the performance of the model with respect to the learning the masked out original text. (Collectively, block **814**).

Attention now turns to discussion of the use of the neural transformer model in an inference system.

Inference Phase

The inference phase of the automated program repair tool relies on a beam search to generate one or more code repair strings corresponding to a buggy code snippet, bug type, and bug-type centroid. The decoder's computation at training time can be parallelized using masked self-attention but during inference, the subtokens are generated one token at a time. The neural transformer model factorizes the probability of the target subtokens in an input sequence into a product of conditional probabilities for each subtoken using the formula: $p(t_1, \dots, t_m | s) = \prod_{i=1}^m p(t_i | t_1, \dots, t_{i-1}, s)$. During inference, the calculation of $\arg \max_s p(t | s)$ is complex and extremely time consuming making the model difficult for real-time applications. Beam search is an approximation algorithm that performs faster.

The beam search uses the probability distribution generated by the neural transformer model to identify the top k subtokens likely to be the next subtoken in a candidate code repair sequence. The beam search expands the search by instantiating new partial sequences using each of the selected subtokens identified by the neural transformer model's probability distribution. The search continues generating new partial sequences from the top k subtokens identified by the output distributions from the neural transformer model until the search ends. The search may end when the end-of-file appears as the most probable next subtoken.

A beam search uses a breadth-first search to build a search tree. The search tree is composed of nodes at one or more inference levels. Each node represents a probability distribution generated by the neural transformer model for the subtokens in the model vocabulary. At each level, only the top k subtokens having the highest probabilities from the output distribution generated by the neural transformer model are expanded to the next inference level. The variable k is preconfigured and referred to as the beam width. Each of the k subtokens is then expanded into a search that

14

updates the current context sequence with the selected subtoken to input into the neural transformer model to generate an additional probability distribution for the next subtoken in a sequence. This process is repeated until the end-of-file token is predicted as being the next likely subtoken candidate.

Turning to FIG. **10A**, there is shown components of the inference phase **1000**. An input sequence **1002** containing a buggy code snippet **1002** and its bug type is provided which is transformed into a corresponding concrete syntax tree that is traversed, by a tokenizer **1004**, to extract tokens and/or subtokens in an ordered sequence. The bug type is added to the ordered sequences at the beginning of the sequence. The ordered sequence of T subtokens is then vectorized into a context vector **1008** using a byte-level byte-pair extraction component **1006**.

The beam search **1010** uses the context vector **1008** to initiate an inference process **1012** using the probability distribution generated from the neural transformer model, $P_0 \dots P_{|V|}$ (block **1012**). If the probability distribution indicates that an end-of-file token is the most likely subtoken to follow in a partial candidate sequence (block **1016**—yes), then the top k candidate sequences are output (block **1018**). Otherwise, the beam search **1010** takes the top k states or subtokens identified from the probability distribution generated by the neural transformer model in the inference process (block **1020**). A new context vector is generated for each of the k states, c_1, \dots, c_k , using the new subtoken in the context vector (blocks **1022A**, **1022B**). The new context vectors are then input into the inference process (blocks **1022A**, **1022B**, **1012**). The beam search **1010** ends when the end-of-file token is selected as the most likely candidate to complete a partial candidate sequence.

FIG. **10B** illustrates an exemplary inference process **1012**. An embedding vector for each subtoken in a sequence **1008** is obtained from the subtoken embedding matrix **1028** and its corresponding positional vector from the positional embedding matrix **1030**. A bug-type centroid **1028** corresponding to the bug type is obtained from the bug-type centroid table **1027**.

The subtoken embedding vector and its corresponding positional embedding vector **1030** are combined to form context tensor **1036**. Context tensor **1036** is input into the context encoder **1042** and concatenated with the edit embedding **1043** output from the edit encoder **1040**. The concatenated embedding **1041** is input into the first decoder block. The context tensor **1034** is input into the first encoder block of the edit encoder **1040** and the output of each encoder block is passed to the next encoder block forming edit embedding **1043**. Edit embedding **1043** is concatenated with the context embedding **1045** and input into each multi-attention layer of each decoder block **204**.

The last decoder block outputs a vector of floating point numbers that is projected by the linear layer **1036** into unnormalized predictions or logits $V_0 \dots V_{|V|}$ **1042**. The logits **1042** are normalized using the softmax function **1044** to generate the softmax prediction **1046** $P_0 \dots P_{|V|}$. Automated Bug Repair Tool

In one aspect, the neural transformer model is used as an automated program repair tool. The automated program repair tool may be part of an integrated development environment (IDE) or part of a web service that analyzes source code for source code bugs and when found, provides a bug fix.

Referring to FIG. **11A**, the automated bug repair tool receives source code to analyze (block **1102**). The source code is analyzed through static analysis, such as the inter-

procedural static analysis tool described above, to identify any potential software bugs and the corresponding bug type (block 1104). The automated bug repair tool performs a beam search using the neural transformer model to predict one or more bug repair candidates (block 1106). The top k candidate bug repairs are returned to the developer (block 1108).

Turning to FIG. 11B, the beam search uses the neural transformer model to generate a probability distribution for the subtoken vocabulary (block 1110). If the probability distribution indicates that the next likely token is the end-of-docstring token, then the beam search is finished (block 1112—yes) and the top k candidate sequences are returned to the user interface (block 1114). Otherwise (block 1112—no), the top k subtokens to complete a partial sequence are selected (block 1116).

Each of the selected subtokens is then input in a respective context vector and has a separate data path through the neural transformer model again. The context vector utilizes the selected subtoken in the current context vector with the last subtoken removed. The new context vector will consist of T subtokens with the selected subtoken t_k added to the beginning of the sequence with the last subtoken removed from the sequence. If the current context vector consists of a subtoken sequence consisting of t_0, t_1, \dots, t_T , then the new context vector will consist of $t_k, t_0, t_1, \dots, t_{T-1}$. (Collectively, block 1118).

Exemplary Operating Environment

Attention now turns to a discussion of an exemplary operating environment. FIG. 12 illustrates an exemplary operating environment 1200 in which one or more computing devices 1202 are used to train the neural transformer model use the neural transformer model for automated program repair. However, it should be noted that the aspects disclosed herein is not constrained to any particular configuration of devices. Computing devices 1202 may be configured as a cloud service that generates the neural transformer model as a service for other code completion systems. It should be noted that the operating environment is not limited to any particular configuration and other configurations are possible.

A computing device 1202 may be any type of electronic device, such as, without limitation, a mobile device, a personal digital assistant, a mobile computing device, a smart phone, a cellular telephone, a handheld computer, a server, a server array or server farm, a web server, a network server, a blade server, an Internet server, a work station, a mini-computer, a mainframe computer, a supercomputer, a network appliance, a web appliance, a distributed computing system, multiprocessor systems, or combination thereof. The operating environment 1200 may be configured in a network environment, a distributed environment, a multiprocessor environment, or a stand-alone computing device having access to remote or local storage devices.

The computing device 1202 may include one or more processors 1206, one or more communication interfaces 1208, one or more storage devices 1210, one or more input/output devices 1212, and one or more memory devices 1214. A processor 1206 may be any commercially available or customized processor and may include dual microprocessors and multi-processor architectures. A communication interface 1208 facilitates wired or wireless communications between the computing device 1202 and other devices. A storage device 1210 may be computer-readable medium that does not contain propagating signals, such as modulated data signals transmitted through a carrier wave. Examples of a storage device 1210 include without limitation RAM,

ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVD), or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage, all of which do not contain propagating signals, such as modulated data signals transmitted through a carrier wave. There may be multiple storage devices 1210 in the computing device 1202. The input/output devices 1214 may include a keyboard, mouse, pen, voice input device, touch input device, display, speakers, printers, etc., and any combination thereof.

A memory device 1214 may be any non-transitory computer-readable storage media that may store executable procedures, applications, and data. The computer-readable storage media does not pertain to propagated signals, such as modulated data signals transmitted through a carrier wave. It may be any type of non-transitory memory device (e.g., random access memory, read-only memory, etc.), magnetic storage, volatile storage, non-volatile storage, optical storage, DVD, CD, floppy disk drive, etc. that does not pertain to propagated signals, such as modulated data signals transmitted through a carrier wave. A memory 1214 may also include one or more external storage devices or remotely located storage devices that do not pertain to propagated signals, such as modulated data signals transmitted through a carrier wave.

The memory device 1214 may contain instructions, components, and data. A component is a software program that performs a specific function and is otherwise known as a module, program, component, and/or application. The memory device 1214 may include an operating system 1218, a pre-training component 1220, a fine-tuning component 1222, a pre-trained neural transformer model 1224, a fine-tuned neural transformer model 1226, source code files 1228, an automated program repair tool 1230, source code repositories 1232, a fine-tuning training dataset 1234, bug edit representation generator 1236, bug-type centroids 1238, and other applications and data 1230.

The computing devices 1202 may be communicatively coupled via a network 1206. The network 1206 may be configured as an ad hoc network, an intranet, an extranet, a virtual private network (VPN), a local area network (LAN), a wireless LAN (WLAN), a wide area network (WAN), a wireless WAN (WWAN), a metropolitan network (MAN), the Internet, a portions of the Public Switched Telephone Network (PSTN), plain old telephone service (POTS) network, a wireless network, a WiFi® network, or any other type of network or combination of networks.

The network 1206 may employ a variety of wired and/or wireless communication protocols and/or technologies. Various generations of different communication protocols and/or technologies that may be employed by a network may include, without limitation, Global System for Mobile Communication (GSM), General Packet Radio Services (GPRS), Enhanced Data GSM Environment (EDGE), Code Division Multiple Access (CDMA), Wideband Code Division Multiple Access (W-CDMA), Code Division Multiple Access 2000, (CDMA-2000), High Speed Downlink Packet Access (HSDPA), Long Term Evolution (LTE), Universal Mobile Telecommunications System (UMTS), Evolution-Data Optimized (Ev-DO), Worldwide Interoperability for Microwave Access (WiMax), Time Division Multiple Access (TDMA), Orthogonal Frequency Division Multiplexing (OFDM), Ultra Wide Band (UWB), Wireless Application Protocol (WAP), User Datagram Protocol (UDP), Transmission Control Protocol/Internet Protocol (TCP/IP), any portion of the Open Systems Interconnection (OSI) model protocols, Session Initiated Protocol/Real-Time Transport

Protocol (SIP/RTP), Short Message Service (SMS), Multimedia Messaging Service (MMS), or any other communication protocols and/or technologies.

CONCLUSION

A system is disclosed comprising one or more processors and a memory. The memory stores one or more programs that are configured to be executed by the one or more processors. The one or more programs include instructions that: obtain a code snippet with a source code bug and an annotated bug type; predict a bug repair for the code snippet from a neural transformer model with attention based on the code snippet and the annotated bug type; and utilize the predicted bug repair to repair the code snippet. In one aspect, the one or more programs include further instructions that: identify the annotated bug type using an interprocedural static analyzer. In an aspect, the one or more programs include further instructions that: pre-train the neural transformer model with an unsupervised training dataset, the unsupervised training dataset including source code snippets.

In an aspect, the one or more programs include further instructions that: fine-tune the pre-trained neural transformer model with a supervised training dataset, the supervised training dataset containing translation tasks, a translation task containing a source code with a bug, a bug type annotation of the bug, and a bug fix for the bug. In an aspect, the one or more programs include further instructions that: generate a bug edit representation for each bug within the supervised training dataset; and compute a bug edit centroid for each bug type based on bug edit representations of each bug type, to be used during inference in place of the edit representation.

In an aspect, the neural transformer model includes one or more encoder blocks and one or more decoder blocks. In an aspect, the one or more programs include further instructions that: utilize the bug-fixing edit representation in at least one or more decoder blocks during training stage, and utilize bug edit centroid in at least one or more decoder blocks during inference stage, the bug edit centroid of a same bug type as the annotated bug type.

A method is disclosed comprising: pre-training a neural transformer model with an unsupervised training dataset, the unsupervised training dataset including a plurality of sequences of source code; fine-tuning the neural transformer model with a supervised training dataset, the supervised training dataset based a triplet including a code snippet with a bug, a code repair for the bug, and an annotated bug type; and applying the neural transformer model to generate a first code repair for a first code snippet having an identified bug and an identified bug type.

In one aspect, the method further comprises: applying a span masking function to each sequence of source code to mask out a subset of subtokens in a sequence; and wherein the neural transformer model learns original subtokens of the sequence. In an aspect, the fine-tuning the neural transformer model with a supervised training dataset further comprises: generating a bug edit embedding representing edits made to correct a bug; and computing a bug type centroid for each bug type from the bug edit embeddings of a particular bug type. In an aspect, the neural transformer model with attention includes one or more encoder blocks coupled to one or more decoder blocks.

In an aspect, fine-tuning the neural transformer model with supervised training dataset further comprises: concatenating the bug-fixing edit embedding with output from a

last encoder block to input to a first decoder block or to encoder-decoder attention block, and concatenating the bug-fixing edit embedding with output embedding at each temporal step. In an aspect, the method, further comprises identifying the annotated bug type through a static analysis of the code snippet. In an aspect, the neural transformer model includes one or more encoder blocks and one or more decoder blocks, wherein an encoder block contains a multi-head attention layer and a feed-forward neural network, wherein a decoder block contains a masked multi-head attention layer, an encoder-decoder multi-head attention layer, and a feed-forward neural network. In an aspect, the annotated bug type includes a null pointer dereference, a memory leak, an immutable cast, empty vector access, or thread safety violation.

A device is disclosed comprising at least one processor and a memory. The at least one processor is configured to train a neural transformer model with attention to learn to translate a source code snippet with a bug and bug type into a code snippet with a repair for the bug by transfer learning, wherein the transfer learning pre-trains the neural transformer model from a plurality of unsupervised training data, the plurality of unsupervised training data including code snippets from a plurality of source code programs, wherein the transfer learning fine-tunes the pre-trained neural transformer model using a plurality of translation tasks, a translation task including a code snippet with a bug, a code snippet with a repair for the bug, and a bug type for the bug; and utilize the neural transformer model to predict a code repair for a second code snippet having a bug and a bug type.

In one aspect, the at least one processor is further configured to: utilize a static code analyzer to identify the bug type of the second code snippet. In an aspect, the at least one processor is further configured to: generate a bug edit representation for each translation task; and compute a bug type centroid for each bug type based on each bug edit representation of a bug type. In an aspect, the neural transformer model includes one or more encoder blocks coupled to one or more decoder blocks, wherein output of a last encoder block is input into each of the decoder blocks. In an aspect, the at least one processor is further configured to: concatenate the output of the last encoder block with a bug type centroid of a bug type of a fine-tuning triplet to a first decoder block.

It may be appreciated that the representative methods described herein do not necessarily have to be executed in the order presented, or in any particular order, unless otherwise indicated. Moreover, various activities described with respect to the methods can be executed in serial or parallel fashion, or any combination of serial and parallel operations.

Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as example forms of implementing the claims.

What is claimed:

1. A method for automated program repair, the method comprising:

training a neural transformer model with attention with a dataset to predict a repaired source code for a given source code snippet and a given bug type;
prior to training the neural transformer model with attention, pre-training the neural transformer model with attention on unsupervised source code snippets;

19

identifying a bug type for a source code bug of a code snippet;
 obtaining a bug-type edit centroid for the bug type of the code snippet, wherein the bug-type edit centroid represents edits to fix the bug type;
 generating a bug repair for the code snippet from the neural transformer model with attention given the bug-type edit centroid, wherein the neural transformer model is trained to predict an encoding for each bug type of the training dataset and transform the encoding of each bug type into the bug-type edit centroid; and
 generating a repaired source code for the code snippet based at least in part on the bug repair.

2. The method of claim 1, further comprising:
 wherein the training dataset comprising a plurality of triplets each comprising a source code snippet with a software bug, a repaired source code, and a bug type of the software bug.

3. The method of claim 1, wherein the neural transformer model with attention includes an encoder block to generate an edit embedding for the bug-type edit centroid.

4. The method of claim 3, further comprising:
 deploying the neural transformer model with attention in an inference system to predict repaired code for a given code snippet having a software bug of a specified bug type, wherein the encoder block receives the bug-type edit centroid of the specified bug type.

5. The method of claim 1, further comprising:
 forming a cluster for each bug type, wherein the cluster includes encodings having a same bug type; and
 computing the bug-type edit centroid for each cluster from the encodings of each bug type in a respective cluster.

6. The method of claim 1, wherein the repaired source code comprises a sequence of source code tokens.

7. A system for automated program repair, the system comprising:
 a processor; and
 a memory including instructions, which when executed by the processor, cause the processor to perform operations comprising:
 training the neural transformer model with attention with a dataset to predict a repaired source code for a given source code snippet and a given bug type;
 prior to training the neural transformer model with attention, pre-training the neural transformer model with attention on unsupervised source code snippets;
 identifying a bug type for a source code bug of a code snippet;
 obtaining a bug-type edit centroid for the bug type of the code snippet, wherein the bug-type edit centroid represents edits to fix the bug type;
 generating a bug repair for the code snippet from the neural transformer model with attention given the bug-type edit centroid, wherein the neural transformer model is trained to predict an encoding for each bug type of the training dataset and transform the encoding of each bug type into the bug-type edit centroid; and
 generating a repaired source code for the code snippet based at least in part on the bug repair.

8. The system of claim 7, wherein the operations further comprise:
 wherein the training dataset comprising a plurality of triplets each comprising a source code snippet with a software bug, a repaired source code, and a bug type of the software bug.

20

9. The system of claim 7, wherein the neural transformer model with attention includes an encoder block to generate an edit embedding for the bug-type edit centroid.

10. The system of claim 9, wherein the operations further comprise:
 deploying the neural transformer model with attention in an inference system to predict repaired code for a given code snippet having a software bug of a specified bug type, wherein the encoder block receives the bug-type edit centroid of the specified bug type.

11. The system of claim 7, wherein the operations further comprise:
 forming a cluster for each bug type, wherein the cluster includes encodings having a same bug type; and
 computing the bug-type edit centroid for each cluster from the encodings of each bug type in a respective cluster.

12. The system of claim 7, wherein the repaired source code comprises a sequence of source code tokens.

13. A non-transitory computer-readable medium including instructions, which when executed by a processor, cause the processor to perform operations comprising:
 training a neural transformer model with attention with a dataset to predict a repaired source code for a given source code snippet and a given bug type;
 prior to training the neural transformer model with attention, pre-training the neural transformer model with attention on unsupervised source code snippets;
 identifying a bug type for a source code bug of a code snippet;
 obtaining a bug-type edit centroid for the bug type of the code snippet, wherein the bug-type edit centroid represents edits to fix the bug type;
 generating a bug repair for the code snippet from the neural transformer model with attention given the bug-type edit centroid, wherein the neural transformer model is trained to predict an encoding for each bug type of the training dataset and transform the encoding of each bug type into the bug-type edit centroid; and
 generating a repaired source code for the code snippet based at least in part on the bug repair.

14. The non-transitory computer-readable medium of claim 13, wherein the operations further comprise:
 wherein the training dataset comprising a plurality of triplets each comprising a source code snippet with a software bug, a repaired source code, and a bug type of the software bug.

15. The non-transitory computer-readable medium of claim 13, wherein the neural transformer model with attention includes an encoder block to generate an edit embedding for the bug-type edit centroid.

16. The non-transitory computer-readable medium of claim 15, wherein the operations further comprise:
 deploying the neural transformer model with attention in an inference system to predict repaired code for a given code snippet having a software bug of a specified bug type, wherein the encoder block receives the bug-type edit centroid of the specified bug type.

17. The non-transitory computer-readable medium of claim 13, wherein the operations further comprise:
 forming a cluster for each bug type, wherein the cluster includes encodings having a same bug type; and
 computing the bug-type edit centroid for each cluster from the encodings of each bug type in a respective cluster.

* * * * *