

(12) **United States Patent**
Frolikov et al.

(10) **Patent No.:** **US 12,393,510 B2**
(45) **Date of Patent:** ***Aug. 19, 2025**

(54) **NAMESPACE MANAGEMENT FOR
MEMORY SUB-SYSTEMS**

(71) Applicant: **Micron Technology, Inc.**, Boise, ID
(US)

(72) Inventors: **Alexei Frolikov**, San Jose, CA (US);
Mark Ish, San Ramon, CA (US)

(73) Assignee: **Micron Technology, Inc.**, Boise, ID
(US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 0 days.

This patent is subject to a terminal dis-
claimer.

(21) Appl. No.: **18/141,870**

(22) Filed: **May 1, 2023**

(65) **Prior Publication Data**

US 2023/0350798 A1 Nov. 2, 2023

Related U.S. Application Data

(63) Continuation of application No. 16/914,939, filed on
Jun. 29, 2020, now Pat. No. 11,675,695.

(51) **Int. Cl.**
G06F 12/02 (2006.01)
G06F 12/0806 (2016.01)
G06F 12/0873 (2016.01)

(52) **U.S. Cl.**
CPC **G06F 12/0246** (2013.01); **G06F 12/0806**
(2013.01); **G06F 12/0873** (2013.01); **G06F**
2212/7201 (2013.01)

(58) **Field of Classification Search**
CPC G06F 12/0246; G06F 12/0806; G06F
12/0873; G06F 2212/7201

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

10,216,630 B1 2/2019 Shang et al.
10,990,526 B1 4/2021 Lam et al.
(Continued)

FOREIGN PATENT DOCUMENTS

CN 110908927 A 3/2020
CN 111344683 A 6/2020

OTHER PUBLICATIONS

Chinese patent office, "China Office Action," issued in connection
with China Patent Application No. 202110702652.X dated Nov. 9,
2024 (20 pages) (9 pages of English Translation and 11 pages of
Original Document).

(Continued)

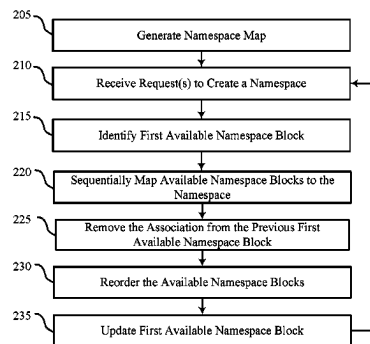
Primary Examiner — Prasith Thammavong

(74) *Attorney, Agent, or Firm* — Holland & Hart LLP

(57) **ABSTRACT**

Methods, systems, and devices for clock domain crossing
queue are described. A memory sub-system can generate a
namespace map having a set of namespace blocks associated
with a memory sub-system. The namespace blocks can
include one or more logical block addresses associated with
the memory sub-system. One namespace block of the set of
namespace blocks can include an indication that can indicate
that the namespace block and each namespace block fol-
lowing the namespace block are available for mapping. The
memory sub-system can receive a request to create a
namespace and sequentially map one or more available
namespace blocks to the namespace according to the order-
ing of the namespace map, including the namespace block
with the indication.

20 Claims, 6 Drawing Sheets



(56)

References Cited

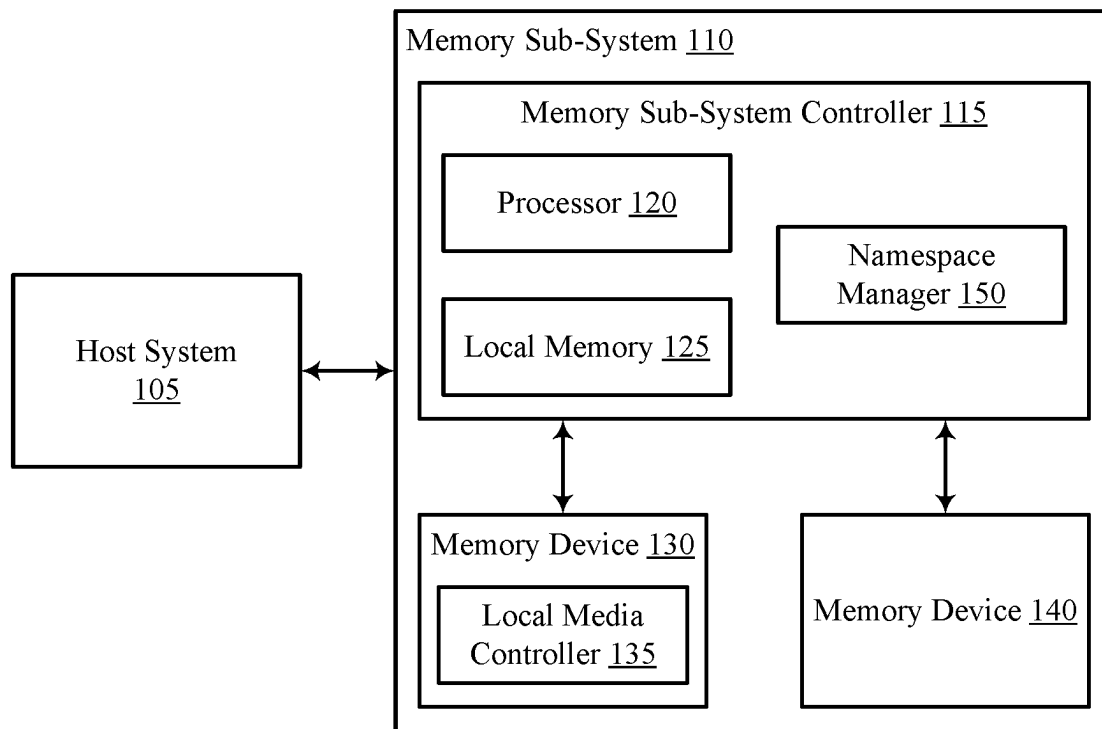
U.S. PATENT DOCUMENTS

11,288,180	B2 *	3/2022	Frolikov	G06F 12/0238
2013/0227201	A1	8/2013	Talagala et al.	
2018/0314441	A1 *	11/2018	Suryanarayana	G06F 3/0683
2018/0335947	A1 *	11/2018	Martineau	G06F 3/0611
2019/0146907	A1 *	5/2019	Frolikov	G06F 12/0246
				711/103
2019/0146927	A1 *	5/2019	Frolikov	G06F 12/0246
				711/103
2020/0089619	A1	3/2020	Hsu et al.	
2020/0409567	A1	12/2020	Bueb et al.	
2021/0191638	A1	6/2021	Miladinovic et al.	
2021/0255949	A1	8/2021	Guda et al.	
2021/0263674	A1	8/2021	Shin et al.	
2023/0315290	A1	10/2023	Frolikov	

OTHER PUBLICATIONS

Zhao, Q., et al., "Research and Implementation of Addressing Method Based on NAND Flash Multichannel Storage System.", Microelectronics & Computer, vol. 35, No. 9, 2018, pp. 32-36.

* cited by examiner



100

FIG. 1

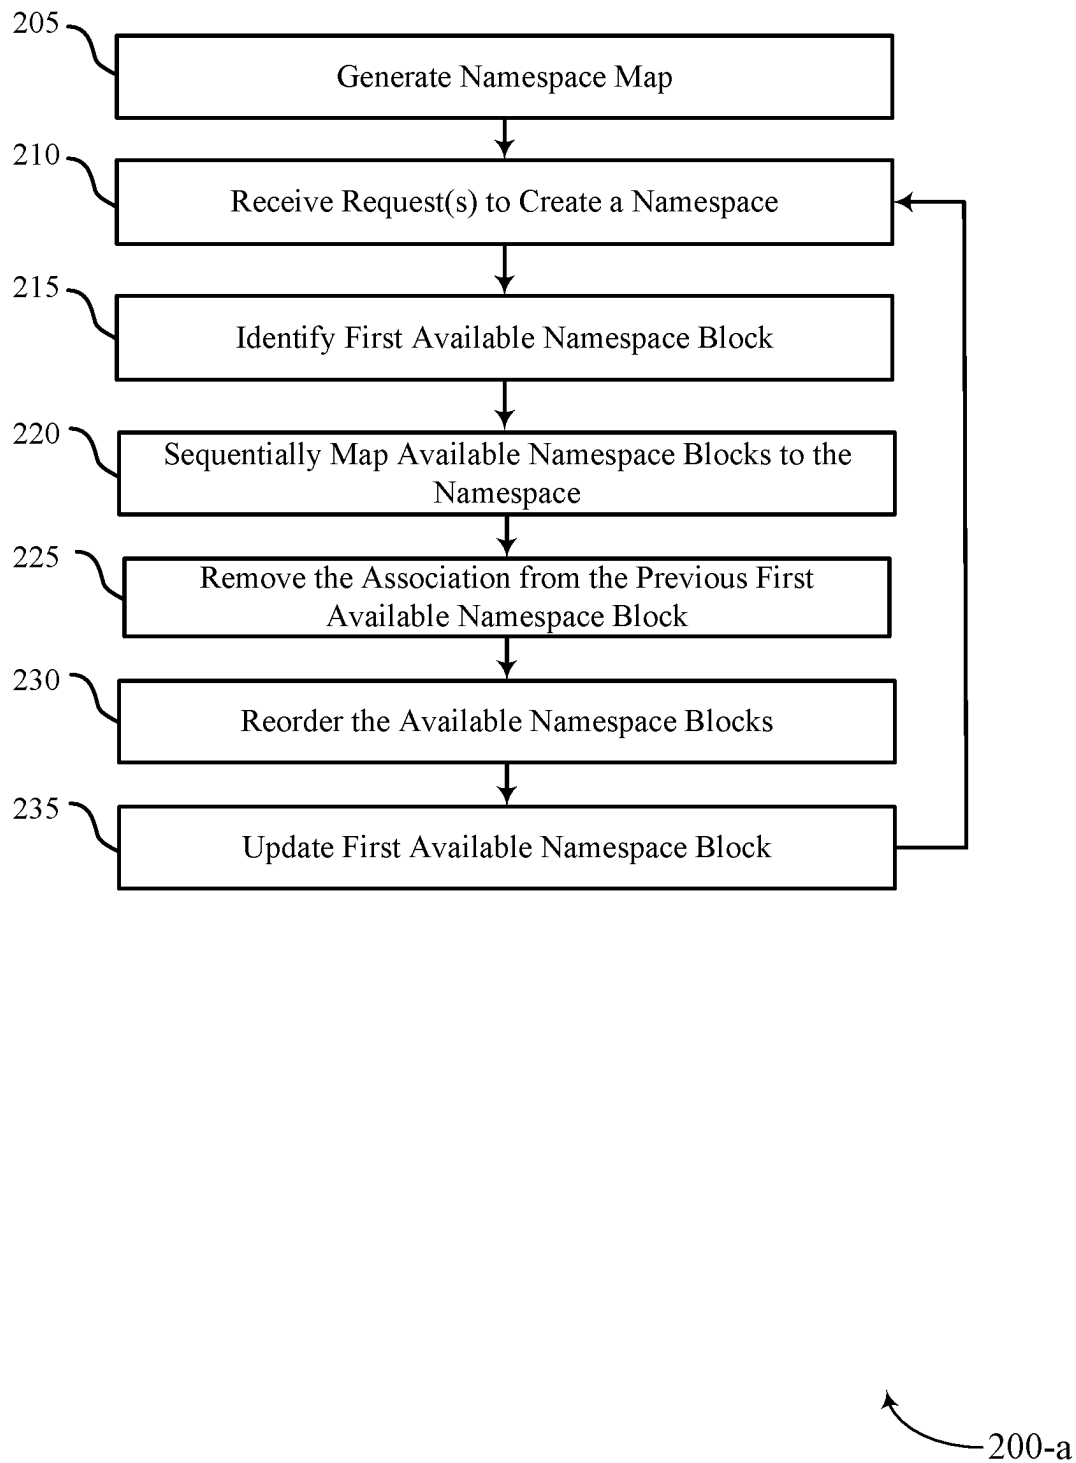


FIG. 2A

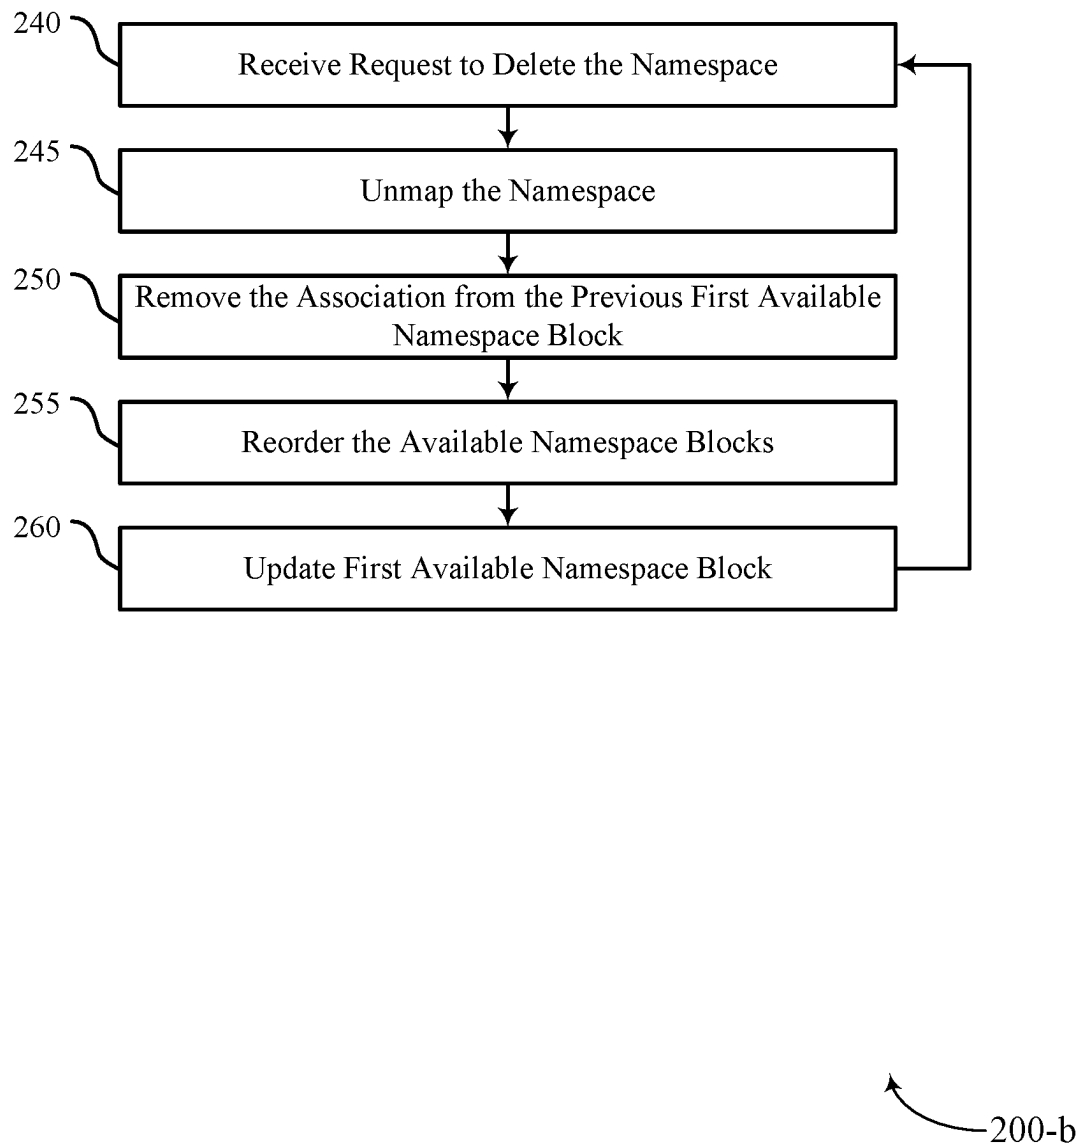
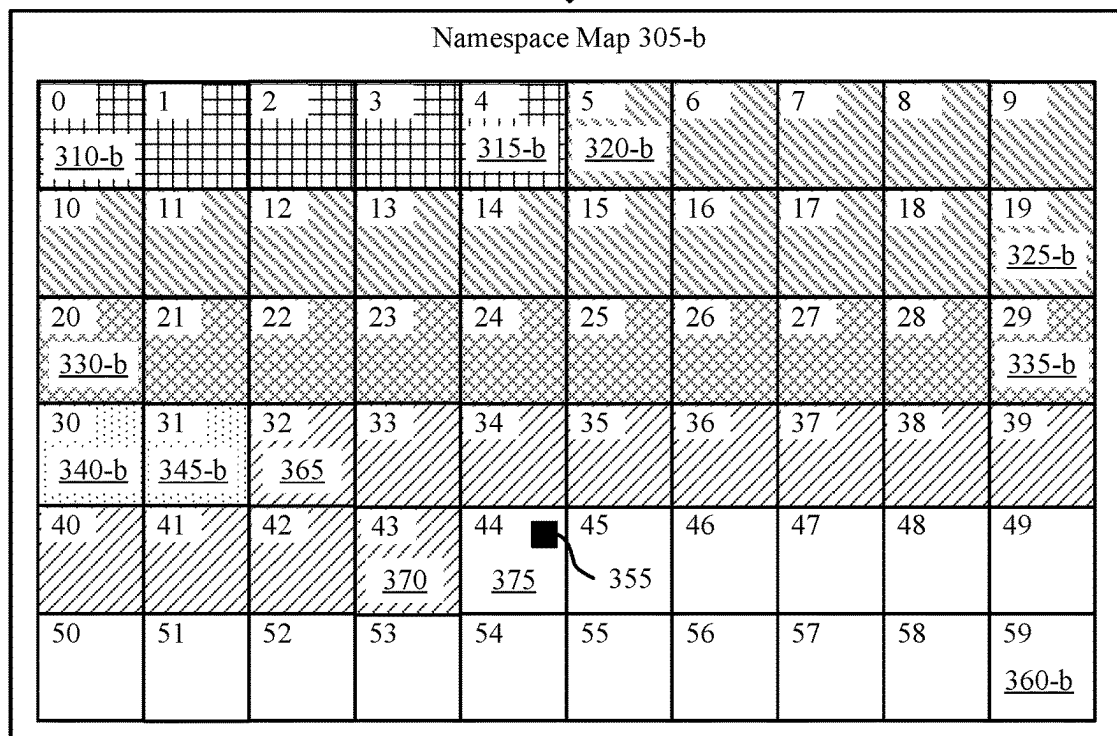

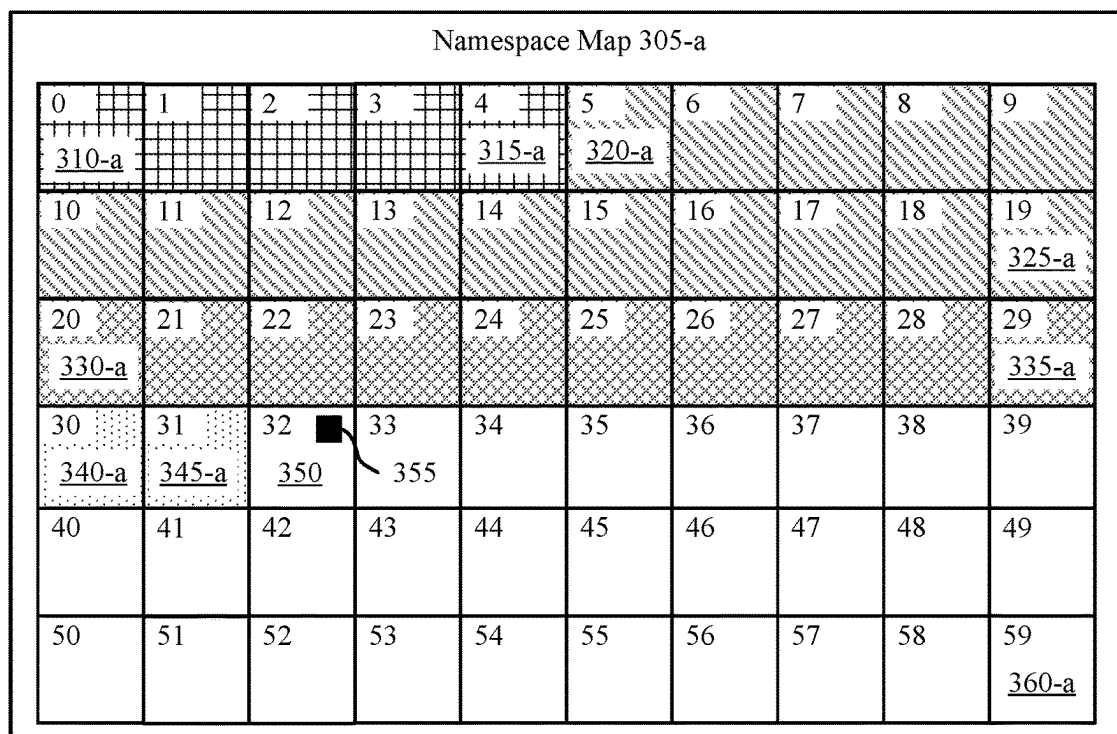
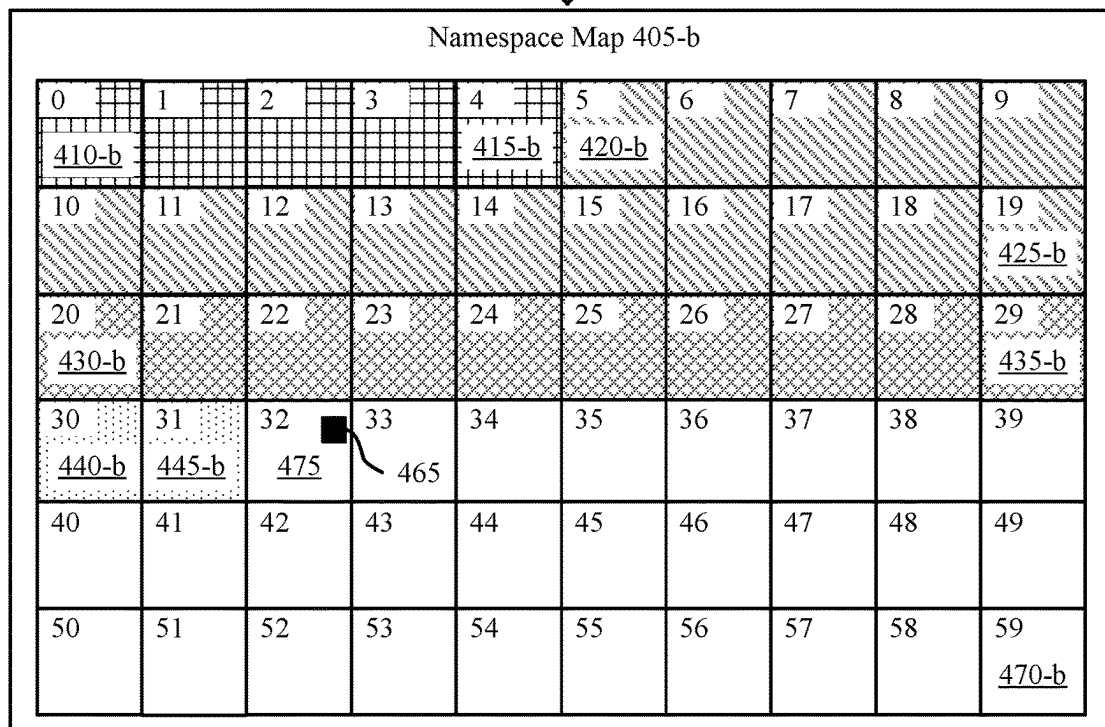




FIG. 2B



300

FIG. 3



400

FIG. 4

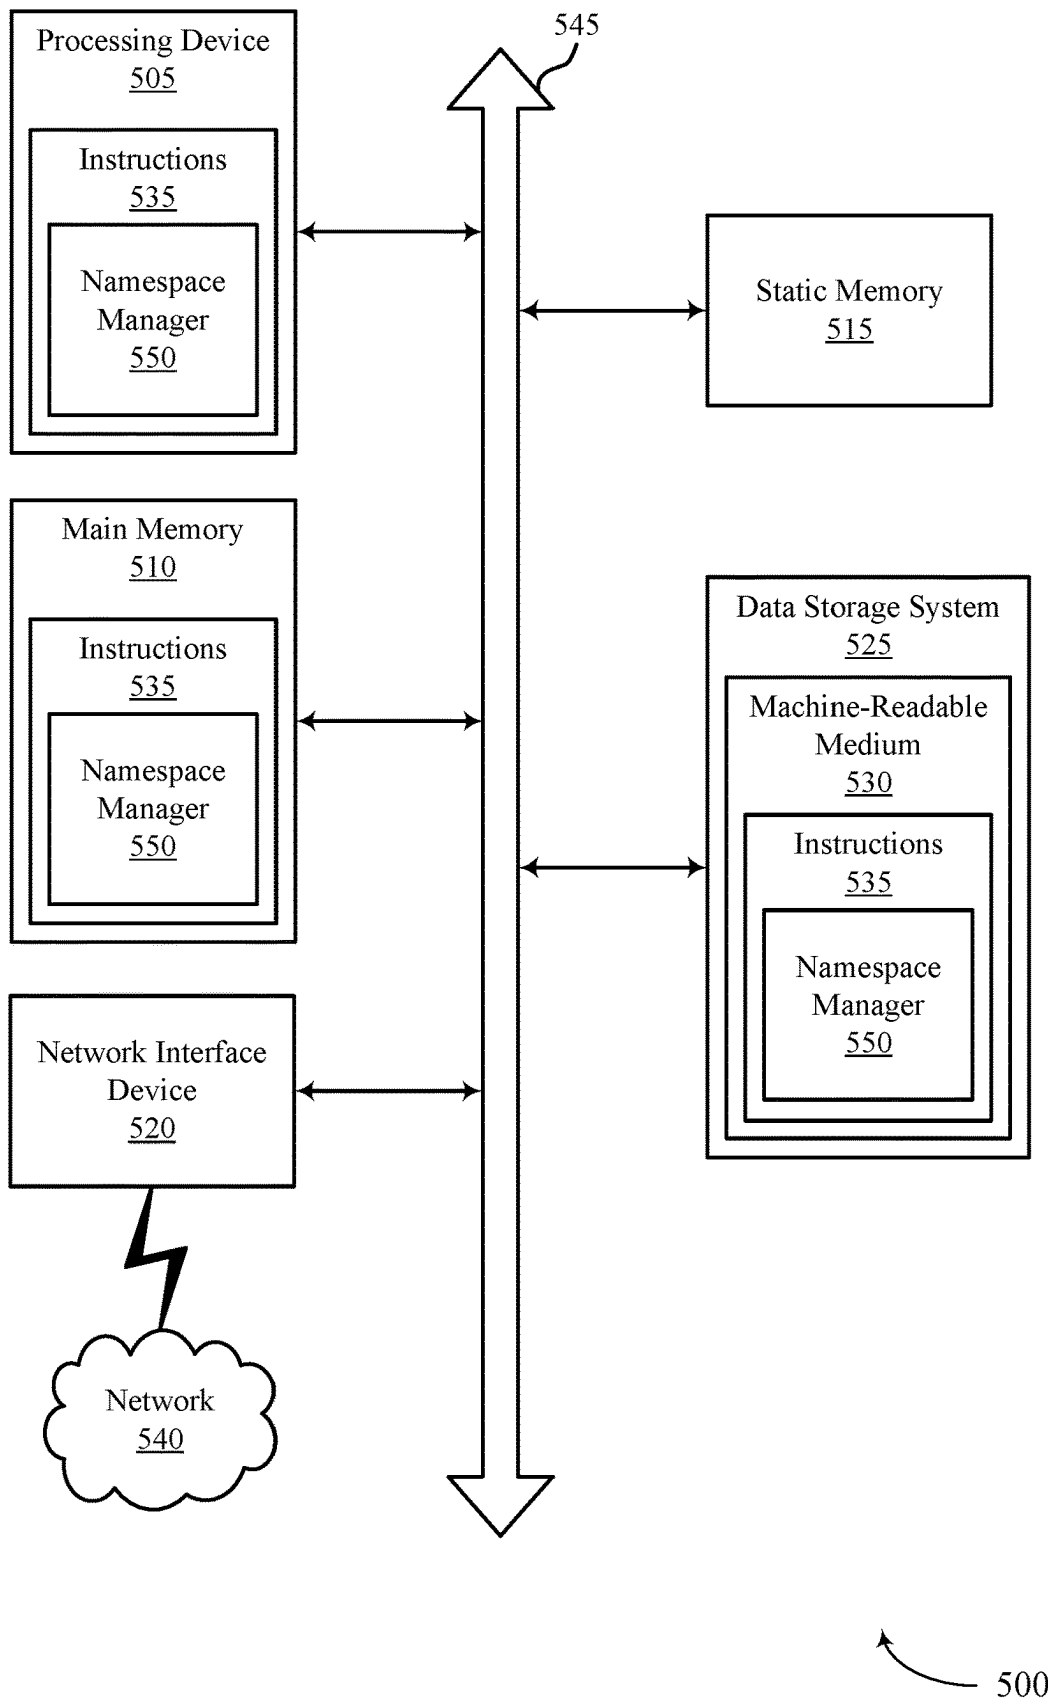


FIG. 5

1

NAMESPACE MANAGEMENT FOR MEMORY SUB-SYSTEMS

CROSS REFERENCE

The present application for patent is a continuation of U.S. patent application Ser. No. 16/914,939 by Frolikov et al., entitled "NAMESPACE MANAGEMENT FOR MEMORY SUB-SYSTEMS," filed Jun. 29, 2020, assigned to the assignee hereof, and is expressly incorporated by reference in its entirety herein.

TECHNICAL FIELD

The following relates generally to a memory sub-system and more specifically to namespace management for memory sub-systems.

BACKGROUND

A memory sub-system can include one or more memory devices that store data. The memory devices can be, for example, non-volatile memory devices and volatile memory devices. In general, a host system can utilize a memory sub-system to store data at the memory devices and to retrieve data from the memory devices.

BRIEF DESCRIPTION OF THE DRAWINGS

The disclosure will be understood more fully from the detailed description given below and from the accompanying drawings of various embodiments of the disclosure. The drawings, however, should not be taken to limit the disclosure to the specific embodiments, but are for explanation and understanding only.

FIG. 1 is an example of a computing environment in accordance with some embodiments of the present disclosure.

FIGS. 2A & 2B are flow diagrams of example methods to manage a namespace in accordance with some embodiments of the present disclosure.

FIG. 3 illustrates an example namespace management operation in accordance with some embodiments of the present disclosure.

FIG. 4 illustrates an example namespace management operation in accordance with some embodiments of the present disclosure.

FIG. 5 is an example machine of a computer system in which embodiments of the present disclosure can operate.

DETAILED DESCRIPTION

Aspects of the present disclosure are directed to sequential prefetching through a linking array. A memory sub-system can be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of storage devices and memory modules are described below in conjunction with FIG. 1. In general, a host system can utilize a memory sub-system that includes one or more memory components (also hereinafter referred to as "memory devices"). The host system can provide data to be stored at the memory sub-system and can request data to be retrieved from the memory sub-system.

A memory device can be a non-volatile memory device. A non-volatile memory device is a package of one or more dice. Each die can consist of one or more planes. For some types of non-volatile memory devices (e.g., negative-and

2

(NAND) devices), each plane consists of a set of physical blocks. Each block consists of a set of pages. Each page consists of a set of memory cells, which store bits of data. For some memory devices, such as NAND devices, blocks are the smallest area than can be erased and pages within the blocks cannot be erased individually. For such devices, erase operations are performed one block at a time.

A page of a block can contain valid data, invalid data, or no data. Invalid data is data that is marked as outdated as a new version of the data is stored on the memory device. Invalid data includes data that was previously written but is no longer associated with a valid logical address, such as a logical address referenced by a host system in a physical to logical (P2L) mapping table. Valid data is the most recent version of such data being stored on the memory device. A memory sub-system can mark data as invalid based on information received, for example, from an operating system. A page that does not contain data includes a page that has been previously erased, and not yet written to.

A memory sub-system can contain multiple memory cells, and each memory cell can be associated with a logical block address (LBA) within firmware of the memory sub-system. The LBAs can be stored in the firmware of the memory sub-system, and the LBAs can be used by the firmware when carrying out access operations such as when reading, writing, erasing, and re-writing data to the memory cells. In some cases, the LBAs can be organized into groups of LBAs. For example, LBAs can be contiguously ordered and organized into groups of LBA blocks. In some cases, the LBA blocks can be referred to as namespace blocks. For example, a memory sub-system can contain 8 terabytes (TB) of memory cells, and the LBAs can be grouped into 2 gigabyte (GB) namespace blocks, where each namespace block contains 2 GB of contiguously ordered LBAs. In this case, 64,000 different namespace blocks are used to organize the LBAs within the firmware (e.g., 8 TB of memory divided by 2 GB per block yields 64,000 LBA blocks).

The memory sub-system can utilize a namespace map to manage the namespace blocks within the firmware. In some cases, the namespace map can be organized into different sections. In the case of Non-Volatile Memory Express (NVMe) protocol, the sections can be referred to as namespaces. The firmware can partition the namespace map into multiple namespaces, and the namespace blocks can be allocated into each namespace. The firmware can allocate all of the namespace blocks to the namespaces until all of the namespace blocks are completely allocated, or the firmware can only allocate some of the namespace blocks to namespaces.

The firmware can utilize a free list when allocating the different namespace blocks into the namespaces. Initially, all namespace blocks can be assigned to the free list such that the free list contains all of the namespace blocks that are available for allocating to a namespace. The firmware can allocate the namespace blocks to the namespaces by reassigning the namespace blocks from the free list to a given namespace (e.g., moved from the free list to the namespace). In some cases, the firmware can continue to allocate the free namespace blocks to the namespaces until there are no longer any free namespace blocks available in the free list.

When utilizing a free list, the size of the free list is configured to be large enough to contain all of the available namespace blocks (e.g., when all of the namespace blocks are available for namespace allocation). Additionally, the namespace map can be configured to be large enough to contain all of the namespace blocks (e.g., when all of the namespace blocks have been allocated to namespaces). In

these cases, there can be redundancy in the amount of firmware storage that is dedicated to managing the namespaces because both the free list and the namespace map are large enough to contain all of the namespace blocks. For example, in the case where 8 TB of memory is available and the namespace block size is 2 GB, both the free list and the namespace map each can contain enough memory to store 64,000 entries of namespace blocks. This redundancy can occupy memory that the firmware could otherwise use to improve efficiency and latency in functions of the memory sub-system.

Aspects of the present disclosure address the above and other deficiencies by providing a memory sub-system that does not include or utilize a separate free list for namespace management. Rather than utilizing and managing both a free list and a namespace map individually, the firmware can utilize a namespace map that specifies information for the namespace blocks, such as whether a given namespace block is available for mapping, without having a separate free list that contains such information. For example, an indicator within the namespace map or stored as a private variable accessible by firmware can be used to differentiate between allocated (e.g., mapped) and unallocated (e.g., unmapped) namespaces rather than by using a free list that includes unmapped LBAs available for namespace mapping and a namespace map having only LBAs mapped to namespaces. Through the use of the namespace map without the use of the free list, the redundancy of the namespace map and free list can be avoided and the firmware can function more efficiently than with the use of the free list. Such techniques can free up memory used by the memory sub-system.

Features of the disclosure are initially described in the context of a computing environment as described with reference to FIG. 1. Features of the disclosure are described in the context of a flow diagram of an example method to manage a namespace without the use of a free list and example namespace management operations without using a free list as described with reference to FIGS. 2-4. These and other features of the disclosure are further illustrated by and described with reference to a computer system that relates to namespace management for memory sub-systems as described with reference to FIG. 5.

FIG. 1 illustrates an example computing system **100** that includes a memory sub-system **110** in accordance with some embodiments of the present disclosure. The memory sub-system **110** can include media, such as one or more non-volatile memory devices (e.g., memory device **130**), one or more volatile memory devices (e.g., memory device **140**), or a combination thereof.

A memory sub-system **110** can be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of a storage device include a solid-state drive (SSD), a flash drive, a universal serial bus (USB) flash drive, an embedded Multi-Media Controller (eMMC) drive, a Universal Flash Storage (UFS) drive, a secure digital (SD) card, and a hard disk drive (HDD). Examples of memory modules include a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), and a non-volatile DIMM (NVDIMM).

The computing system **100** can be a computing device such as a desktop computer, laptop computer, network server, mobile device, a vehicle (e.g., airplane, drone, train, automobile, or other conveyance), Internet of Things (IoT) enabled device, embedded computer (e.g., one included in a

vehicle, industrial equipment, or a networked commercial device), or such computing device that includes memory and a processing device.

The computing system **100** can include a host system **105** that is coupled with one or more memory sub-systems **110**. In some examples, the host system **105** is coupled with different types of memory sub-systems **110**. FIG. 1 illustrates one example of a host system **105** coupled with one memory sub-system **110**. As used herein, “coupled to” or “coupled with” generally refers to a connection between components, which can be an indirect communicative connection or direct communicative connection (e.g., without intervening components), whether wired or wireless, including connections such as electrical, optical, magnetic, etc.

The host system **105** can include a processor chipset and a software stack executed by the processor chipset. The processor chipset can include one or more cores, one or more caches, a memory controller (e.g., NVDIMM controller), and a storage protocol controller (e.g., PCIe controller, SATA controller). The host system **105** uses the memory sub-system **110**, for example, to write data to the memory sub-system **110** and read data from the memory sub-system **110**.

The host system **105** can be coupled to the memory sub-system **110** using a physical host interface. Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, USB interface, Fiber Channel, Small Computer System Interface (SCSI), Serial Attached SCSI (SAS), a double data rate (DDR) memory bus, a dual in-line memory module (DIMM) interface (e.g., DIMM socket interface that supports Double Data Rate (DDR)), Open NAND Flash Interface (ONFI), Double Data Rate (DDR), Low Power Double Data Rate (LPDDR), or any other interface. The physical host interface can be used to transmit data between the host system **105** and the memory sub-system **110**. The host system **105** can further utilize a non-volatile memory Express (NVMe) interface to access the memory components (e.g., memory devices **130**) when the memory sub-system **110** is coupled with the host system **105** by the PCIe interface. The physical host interface can provide an interface for passing control, address, data, and other signals between the memory sub-system **110** and the host system **105**.

The memory devices **130, 140** can include any combination of the different types of non-volatile memory devices and/or volatile memory devices. The volatile memory devices (e.g., memory device **140**) can be, but are not limited to, random access memory (RAM), such as dynamic RAM (DRAM) and synchronous DRAM (SDRAM).

Some examples of non-volatile memory devices (e.g., memory device **130**) includes a NAND type flash memory and write-in-place memory, such as a three-dimensional cross-point (“3D cross-point”) memory device, which is a cross-point array of non-volatile memory cells. A cross-point array of non-volatile memory can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many flash-based memories, cross-point non-volatile memory can perform a write in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased. NAND type flash memory includes, for example, two-dimensional NAND (2D NAND) and three-dimensional NAND (3D NAND).

Each of the memory devices **130** can include one or more arrays of memory cells. One type of memory cell, for example, single level cells (SLC) can store one bit per cell. Other types of memory cells, such as multi-level cells (MLCs), triple level cells (TLCs), and quad-level cells (QLCs), can store multiple bits per cell. In some embodiments, each of the memory devices **130** can include one or more arrays of memory cells such as SLCs, MLCs, TLCs, QLCs, or any combination of such. In some embodiments, a particular memory device can include an SLC portion, and an MLC portion, a TLC portion, or a QLC portion of memory cells. The memory cells of the memory devices **130** can be grouped as pages that can refer to a logical unit of the memory device used to store data. With some types of memory (e.g., NAND), pages can be grouped to form blocks.

Although non-volatile memory devices such as NAND type flash memory are described, the memory device **130** can be based on any other type of non-volatile memory, such as read-only memory (ROM), phase change memory (PCM), self-selecting memory, other chalcogenide based memories, ferroelectric transistor random-access memory (FeTRAM), ferroelectric RAM (FeRAM), magneto RAM (MRAM), negative-or (NOR) flash memory, Spin Transfer Torque (STT)-MRAM, conductive bridging RAM (CBRAM), resistive random access memory (RRAM), oxide based RRAM (OxRAM), and electrically erasable programmable ROM (EEPROM).

The memory sub-system controller **115** (or controller **115** for simplicity) can communicate with the memory devices **130** to perform operations such as reading data, writing data, or erasing data at the memory devices **130** and other such operations. The memory sub-system controller **115** can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, or a combination thereof. The hardware can include digital circuitry with dedicated (i.e., hard-coded) logic to perform the operations described herein. The memory sub-system controller **115** can be a microcontroller, special purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), a digital signal processor (DSP)), or other suitable processor.

The memory sub-system controller **115** can include a processor **120** (e.g., a processing device) configured to execute instructions stored in a local memory **125**. In the illustrated example, the local memory **125** of the memory sub-system controller **115** includes an embedded memory configured to store instructions for performing various processes, operations, logic flows, and routines that control operation of the memory sub-system **110**, including handling communications between the memory sub-system **110** and the host system **105**.

In some examples, the local memory **125** can include memory registers storing memory pointers, fetched data, etc. The local memory **125** can also include read-only memory (ROM) for storing micro-code. While the example memory sub-system **110** in FIG. 1 has been illustrated as including the memory sub-system controller **115**, in another example of the present disclosure, a memory sub-system **110** does not include a memory sub-system controller **115**, and can instead rely upon external control (e.g., provided by an external host, or by a processor or controller separate from the memory sub-system).

In general, the memory sub-system controller **115** can receive commands or operations from the host system **105** and can convert the commands or operations into instructions or appropriate commands to achieve the desired access

to the memory devices **130**. The memory sub-system controller **115** can be responsible for other operations such as wear leveling operations, garbage collection operations, error detection and error-correcting code (ECC) operations, encryption operations, caching operations, and address translations between a logical address (e.g., logical block address (LBA), namespace) and a physical address (e.g., physical block address) that are associated with the memory devices **130**. The memory sub-system controller **115** can further include host interface circuitry to communicate with the host system **105** via the physical host interface. The host interface circuitry can convert the commands received from the host system into command instructions to access the memory devices **130** as well as convert responses associated with the memory devices **130** into information for the host system **105**.

The memory sub-system **110** can also include additional circuitry or components that are not illustrated. In some examples, the memory sub-system **110** can include a cache or buffer (e.g., DRAM) and address circuitry (e.g., a row decoder and a column decoder) that can receive an address from the memory sub-system controller **115** and decode the address to access the memory devices **130**.

In some examples, the memory devices **130** include local media controllers **135** that operate in conjunction with memory sub-system controller **115** to execute operations on one or more memory cells of the memory devices **130**. An external controller (e.g., memory sub-system controller **115**) can externally manage the memory device **130** (e.g., perform media management operations on the memory device **130**). In some embodiments, a memory device **130** is a managed memory device, which is a raw memory device combined with a local controller (e.g., local controller **135**) for media management within the same memory device package. An example of a managed memory device is a managed NAND (MNAND) device.

The memory sub-system **110** includes a namespace manager **150** that can manage the namespaces and LBA blocks available or mapped to the namespaces (which can be referred to as namespace blocks) within the firmware of the memory sub-system **110**. The namespace manager **150** can contain all of the namespace blocks associated with the memory sub-system **110**, and can manage the namespace blocks through mapping of the namespace blocks to namespaces within namespace manager **150**. In some cases, an indicator can be used to determine which namespace blocks are allocated to the namespaces, and the indicator can allow for the mapping of namespace blocks to namespaces without the use of a free list.

In some examples, the memory sub-system controller **115** includes at least a portion of the namespace manager **150**. For example, the memory sub-system controller **115** can include a processor **120** (e.g., a processing device) configured to execute instructions stored in local memory **125** for performing the operations described herein. In some examples, the namespace manager **150** is part of the host system **105**, an application, or an operating system.

The namespace manager **150** can map the namespace blocks to the namespaces without using a free list separate from the namespace manager **150**. In some cases, the namespace map can contain an indicator associated with a namespace block within the namespace. The indicator can indicate the division between mapped namespace blocks and namespace blocks available for mapping (e.g., unmapped namespace blocks). In this case, when a new namespace is mapped by the firmware, the firmware can sequentially map the new namespace with the available namespace blocks,

starting with the namespace block associated with the indicator (e.g., in cases where the indicator designates the first namespace block available for mapping). After mapping the new namespace, the indicator can be disassociated with the previous first available namespace block and associated with a new available namespace block, which can be the first available namespace block for namespace mapping. In some cases, the indicator can be a private variable stored in memory accessible by firmware and can be or can convey the index of the first available namespace block. In the case where a namespace is deleted, the indicator can be disassociated from the previously first available namespace block, and after reordering of the navigable namespace blocks, associated with the new (i.e., updated as a result of the namespace mapping) first available namespace block. In these cases, the namespace manager **150** can be managed (e.g., namespaces can be added and deleted) without the use of a free list. Further details with regards to the operations of the namespace manager **150** are described below.

FIG. 2A is a flow diagram of an example method **200-a** to manage namespaces of a memory sub-system in accordance with some embodiments of the present disclosure. The method **200-a** can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **200-a** is performed by the namespace manager **150** of FIG. 1. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated examples should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various examples. Thus, not all processes are required in every example. Other process flows are possible.

At operation **205**, the processing device generates a namespace map. In some embodiments, the namespace map is generated one time for the lifetime of the memory sub-system. For example, one or more default namespace maps can be generated at the time of manufacturing of the memory sub-system or at time of first boot-up. In other cases, the namespace map is generated more than once over the lifetime of the memory sub-system. For example, if the memory sub-system undergoes formatting during its lifetime (e.g., a low level format to reset the drive to a manufacturing state). The namespace map can contain enough entries to organize each namespace block within the namespace map. As discussed above, each namespace block can contain groups of contiguously ordered LBAs (e.g., sequentially ordered LBAs) associated with the memory cells within the memory sub-system. In this case, the size of the namespace map can be based upon the number of namespace blocks in the memory sub-system and therefore on the number of LBAs, which can be grouped into each namespace block. For example, the firmware can organize the LBAs into 2 GB namespace blocks. In this case, the namespace map can contain enough entries for each 2 GB namespace block, and the namespace map size can be based upon this amount of entries. In some cases, the firmware can change the namespace block size, and the size of the namespace map can accommodate the change. For example, the firmware can decrease the size of the namespace block to 1 GB. In this case, the namespace map can double in the amount of entries to accommodate twice as many namespace blocks.

At operation **210**, the processing device receives a request to create a namespace. The request can be a host request for namespace management received from a host system, such as host system **105** as discussed with reference to FIG. 1. In some cases, the request can be received upon bootup of the memory sub-system or when a new namespace is requested from a host system **105**. In some embodiments, the request is an opcode, such as a namespace management opcode. Within the opcode, a select instruction specifies a namespace create, which the processing device identifies and uses to generate the namespace. In some cases, the request can include the size of the namespace. The size of the namespace can determine the number of namespace blocks assigned to the namespace. For example, the memory sub-system can receive a request to create a 24 GB namespace. In the case where the namespace block size is set at 2 GB, 12 namespace blocks can be allocated to the namespace (e.g., 12 namespace blocks at 2 GB each yields a 24 GB namespace).

In some cases, the namespace map can be organized according to an index. For example, an index number can be associated with each entry in the namespace map to which a namespace block is allocated. In this case, each namespace can therefore be associated with a group of contiguous index numbers. Namespace blocks can then be allocated to the namespace, and the namespace is associated with the index numbers. In this case, the index numbers can be used to manage the allocation of the namespace blocks to the namespace. In other cases, a namespace map may not be ordered according to index and instead may be a random order.

At operation **215**, the processing device identifies the first available namespace block. For example, an indicator, such as a set of bits of a variable, can be associated with a namespace block of the namespace map that is available for namespace allocation. In some cases, the namespace block can be the first available namespace block for namespace allocation (e.g., the first namespace block that is not allocated to a namespace (e.g., according to an increasing order of namespace indices)). In some embodiments, the variable can be a private variable associated with the namespace map and accessible by the firmware of the memory sub-system (e.g., the variable can be stored on memory of the memory sub-system that is accessible by the firmware). The firmware can then reference this variable when determining available namespace blocks to map to the namespace. The indicator can be used to specify a division in the namespace map between available and unavailable namespace blocks. For example, all of the namespace blocks ordered before (e.g., with index numbers less than) the indicator can be the unavailable namespace blocks, and all of the namespace blocks ordered after (e.g., with index numbers greater than) the indicator can be the available namespace blocks.

At operation **220**, the processing device maps available namespace blocks. For example, available namespace blocks can be sequentially mapped (e.g., allocated) to the namespace. As discussed above, the first available namespace block can be allocated to the namespace. Additional namespace blocks can be allocated to the namespace until the size of the namespace is fully allocated. For example, a 24 GB namespace can be requested, and based upon a namespace block size of 2 GB, 12 namespace blocks can be allocated to fill the namespace. In this case, the first available namespace block is mapped to the namespace, as well as 11 additional namespace blocks. The additional namespace blocks mapped to the namespace can sequentially follow the first available namespace block

according to the index of the namespace block. For example, in the case where the index number of the first available namespace is 32, namespace block 33 through 43 are also mapped to the namespace.

At operation **225**, processing device removes the association of the variable or indicator with the first available namespace block of operation **215** (i.e., the processing device disassociates the indicator from the previous first available namespace block). The indication is disassociated with the namespace block since the namespace block is no longer the first available namespace block, nor a namespace block available for namespace allocation. In this case, the index number associated with the block is also no longer associated with the first available namespace block for namespace allocation.

At operation **230**, the process device reorders the available namespace blocks. As discussed in operation **220**, previously available namespace blocks have been allocated to the namespace. In this case, the index numbers associated with the available and unavailable namespace blocks can no longer correspond to namespace blocks that are available for namespace allocation. In this case, the firmware can reorder the available namespace blocks by assigning the newly available namespace blocks to index numbers corresponding to the available namespace blocks. In some cases, this can reorder all of the available namespace blocks to new index numbers in the namespace map. For example, in the case where the new namespace has been mapped to occupy index numbers 32-43, all of the available namespace blocks (e.g., namespace blocks associated with index numbers greater than 43) can be reordered.

At operation **235**, the processing device updates the first available namespace block that is available for namespace mapping. As discussed previously, at operation **225**, the association of the variable indicating the first available namespace block can be removed from the previously first available namespace, and at operation **230**, the available namespace blocks can be reordered. In this case, the variable or indicator can be written to storage or memory of the memory sub-system that is accessible by firmware and can be associated with or updated with a set of bits that corresponds to an index of a new (e.g., second) namespace block that is available for namespace allocation. The new namespace block can be the now-first available namespace block for namespace allocation (e.g., the first namespace block that is not allocated to a namespace after the namespace has been allocated). The firmware can then allocate this first available namespace block and additional available namespace blocks to another namespace.

In some cases, operations **210-235** can be performed reiteratively to map multiple namespaces in a case where multiple requests for creating namespaces is received at the memory sub-system. For example, a request to create an additional namespace can be received by the memory sub-system (e.g., similar to operation **210**). The memory sub-system can then sequentially allocate multiple available namespace blocks to the namespace (e.g., similar to operation **220**) based on the variable that indicates the first available namespace block for mapping, remove the association from the now-mapped namespace block (e.g., similar to operation **225**), reorder the available namespace blocks (e.g., similar to operation **230**), and associate another namespace block that is available for namespace mapping (e.g., similar to operation **235**) with the indicator or variable. These steps can be repeated in order to map multiple namespaces in the namespace map.

FIG. **2B** is a flow diagram of an example method **200-b** to manage namespaces of a memory sub-system in accordance with some embodiments of the present disclosure. The method **200-b** can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **200-b** is performed by the namespace manager **150** of FIG. **1**. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated examples should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various examples. Thus, not all processes are required in every example. Other process flows are possible.

In some cases, a previously mapped namespace can be deleted from a namespace map. Deletion of a namespace can occur after one or more namespaces have been allocated in the namespace map. For example, at operation **240**, the memory sub-system can receive a request to delete a namespace that has been mapped to the namespace map. The mapping of the namespace can have occurred during operations **210-235** (e.g., mapping a first namespace) or operations similar to operations **210-235** (e.g., mapping a subsequent namespace) as described in FIG. **2A**.

The request can be a host request for namespace management received from a host system, such as host system **105** as discussed with reference to FIG. **1**. In some embodiments, the request is an opcode, such as a namespace management opcode. Within the opcode, a select instruction specifies a namespace delete, which the processing device identifies and uses to delete the namespace. In some cases, the request can include the size of the namespace. The size of the namespace can indicate the number of namespace blocks assigned to the namespace. For example, the memory sub-system can receive a request to delete a 24 GB namespace. In the case where the namespace block size is set at 2 GB, 12 namespace blocks are allocated to the namespace (e.g., 12 namespace blocks at 2 GB each yields a 24 GB namespace).

At operation **245**, the processing device unmaps the previously allocated namespace. Unmapping of the namespace can include disassociating each mapped namespace block from the namespace and associating the now disassociated namespace blocks with namespace blocks available for mapping. In this example, the unmapped namespace blocks can be associated with index numbers that are available for mapping. For example, in the case where a namespace includes namespace blocks associated with index numbers 32 through 43 (e.g., namespace blocks 32 through 43 are mapped to a namespace), and the namespace is deleted, the index numbers previously associated with the namespace (e.g., index numbers 32 through 43) can then be associated with free namespace blocks available for namespace mapping.

At operation **250**, the processing device disassociates the indication (e.g., the variable) that was previously associated with the first available namespace block. In some cases, the indication is removed from association with this namespace block because the namespace block is no longer the first available namespace block for namespace mapping. For example, in the case that the namespace is deleted that was associated with index numbers 32 through 43, the index number associated with the first available namespace block

can change from 44 (e.g., the first available namespace block before namespace deletion) to an index number of 32 (e.g., the new first available namespace block index number after namespace deletion). In this case, the indication previously indicated the namespace block associated with index 44 as can be removed since the namespace block associated with index number 32 is now the first available namespace block for namespace mapping. In some cases, however, after unmapping the namespace, the first available namespace block for mapping is the same as before mapping the namespace and removing the association of the indicator with the first available namespace block is not performed.

At operation 255, the processing device reorders the available namespace blocks in the namespace map. Similar to operation 235, previously mapped namespace blocks can have been unmapped from the namespace. In this case, the index numbers associated with the available and unavailable namespace blocks can no longer correspond to namespace blocks that are available for namespace allocation. In this case, the firmware can reorder the available namespace blocks by assigning the unmapped namespace blocks to index numbers corresponding to available namespace blocks. In some cases, this can reorder all of the available namespace blocks to new index numbers in the namespace map. For example, after the deletion of the namespace previously associated with index numbers 32 through 43, all of the now available namespace blocks (e.g., namespace blocks associated with index numbers greater than 31) can be reordered.

At operation 260, the processing device updates the first available namespace block. As discussed previously, at operation 250, the indication can be disassociated from the previously first-available namespace block, and at operation 255, the namespace map can be reordered. After deleting the namespace and reordering of the available namespace blocks, the indication can be associated with a different namespace block. This namespace block can be the new first-available namespace block in the namespace map. For example, as discussed previously, the namespace associated with index numbers 32 through 43 can be deleted, and the available namespace blocks can be reordered, such that the namespace associated with index number 32 is now the first available namespace block. In this case, the indication is updated to be associated with the namespace block associated with index number 32.

In some cases, operations similar to operations 240-260 can be performed reiteratively to delete (e.g., unmap) multiple namespaces. For example, a request to delete a second namespace can be received by the memory sub-system (e.g., similar to operation 240). The memory sub-system can then unmap the namespace blocks from the namespace (e.g., similar to operation 245), remove the association of the indicator with the previously first available namespace block (e.g., similar to operation 250), reorder the available namespace blocks (e.g., similar to operation 255), and update the indication of the first available namespace block to a new namespace block that is available for namespace mapping (e.g., similar to operation 260). These steps can be repeated in order to unmap multiple namespaces in the namespace map.

FIG. 3 illustrates an example namespace management operation 300 where a namespace has been deleted without the use of a free list. In some cases, the namespace management operation 300 can include some of the processing steps of method 200-a, as performed by namespace manager 150 as illustrated in FIG. 1.

The namespace management operation 300 can include the addition of a namespace to a namespace map. Namespace map 305-a and namespace map 305-b illustrate two examples of namespace maps showing the change in the namespace map before namespace creation, as indicated by namespace map 305-a, and after namespace creation, as indicated by namespace map 305-b.

Namespace map 305-a illustrates an example namespace map before the mapping of a new namespace. Namespace map 305-a can include four mapped namespaces. The first namespace can include five namespace blocks beginning at namespace block 310-a (e.g., index number 0) and including all sequentially ordered namespace blocks until, and including, namespace block 315-a (e.g., index number 4). The second namespace can include fifteen namespace blocks beginning at namespace block 320-a (e.g., index number 5) and including all sequentially ordered namespace blocks until, and including, namespace block 325-a (e.g., index number 19). The third namespace can include 10 namespace blocks beginning at namespace block 330-a (e.g., associated with index number 20) and including all sequentially ordered namespace blocks until, and including, namespace block 335-a (e.g., associated with index number 29). The fourth namespace can include two namespace blocks beginning at namespace block 340-a (e.g., associated with index number 30) and include all sequentially ordered namespace blocks until, and including, namespace block 345-a (e.g., associated with index number 31). In this case, namespace block 350 can be the first available unmapped namespace block (e.g., associated with index number 32), and all namespace block sequentially ordered after namespace block 350, until, and including namespace block 360-a (e.g., associated with index number 59) can be namespace blocks available for namespace mapping.

In the case of namespace map 305-a, an indication 355 can be a variable (e.g., a private variable written to storage accessible by firmware) and can be associated with namespace block 350 (e.g., as described with reference to operation 215 of method 200-a). In some cases, the indication 355 is a set of bits representative of the index of the namespace block with which the indication 355 is associated. The indication 355 indicates that namespace block 350 is the first available namespace block for namespace mapping. The indication 355 can indicate that all namespace blocks associated with index numbers greater than namespace block 350 are available for namespace mapping (e.g., associated with index numbers greater than or equal to 32). Although namespace blocks with index numbers between 32 and 59 are shown in namespace map 305-a in increasing consecutive order, such an order is non-limiting and any ordering (e.g., a random order) of the namespace blocks may be considered without departing from the scope of the disclosure.

As discussed regarding operation 205 of method 200-a, the memory sub-system can receive a request to create a new namespace. The request can include the size of the namespace, and the memory sub-system can determine the number of namespace blocks to allocate to the namespace. For example, a request can be received for a 24 GB namespace. In the case where each namespace block includes 2 GB of contiguously ordered LBAs, twelve namespace blocks would be allocated to the requested namespace. In this case twelve available namespace blocks would be mapped to the requested namespace. As discussed with reference to operation 220 of method 200-a, the twelve namespace blocks can be mapped to the new namespace according to the ordering of namespace map 305-a.

Namespace map **305-b** illustrates an example namespace map after the mapping of a new namespace. As with namespace map **305-a**, the first four namespaces can have already been mapped and include namespaces associated with index numbers 0 through 31. The newly mapped fifth namespace can include twelve namespace blocks beginning at namespace block **365** (e.g., associated with index number 32) and include all sequentially ordered namespace blocks until, and including, namespace block **370** (e.g., associated with index number 29). As discussed with reference to operation **225** of method **200-a**, after the mapping of the new namespace to namespace map **305-b**, the indication **355** can be disassociated from the previously first-available namespace block. As illustrated in namespace map **305-a**, the indication **355** can have previously been associated with namespace block **350** indicating that namespace block **350** was the first available namespace block. Upon mapping of the new namespace (the fifth namespace), namespace block **350** is now associated with the new namespace and is no longer available for namespace mapping.

In response to the allocation of the available namespace blocks to the new namespace (e.g., namespace blocks associated with index numbers 32 through 43), and as discussed with reference to operation **230** of method **200-a**, the available namespace blocks can be reordered. As discussed with reference to operation **235** of method **200-a**, after reordering of the available namespace blocks, the indication **355** can be updated to be associated with a new first available namespace block. For example, after reordering of the available namespace blocks (e.g., as associated with index numbers 44 through 59), the new first available namespace block can be namespace block **375** (e.g., associated with index number 44). In this case, indication **355** can be associated with namespace block **375**. The indication **355** can indicate that namespace block **375** is now the first available namespace block for namespace mapping, and in the event that a second new namespace is mapped to the memory sub-system, namespace block **375** can be used to indicate namespace blocks available for mapping.

FIG. 4 illustrates an example namespace management operation **400** where a namespace has been deleted without the use of a free list. In some cases, the namespace management operation **400** can include some of the processing steps of method **200-b**, as performed by namespace manager **150** as illustrated in FIG. 1.

The namespace management operation **400** can include the deletion of a namespace from a namespace map. Namespace map **405-a** and namespace map **405-b** illustrate two examples of a namespace map showing the change in the namespace map before namespace deletion, as indicated by names map **405-a**, and after namespace deletion, as indicated by namespace map **405-b**.

Namespace map **405-a** illustrates an example namespace map before the deletion of a namespace from the namespace map. Namespace map **405-a** can include five previously mapped namespaces. The first namespace can include five namespace blocks beginning at namespace block **410-a** (e.g., index number 0) and including all sequentially ordered namespace blocks until, and including, namespace block **415-a** (e.g., index number 4). The second namespace can include fifteen namespace blocks beginning at namespace block **420-a** (e.g., index number 5) and including all sequentially ordered namespace blocks until, and including, namespace block **425-a** (e.g., index number 19). The third namespace can include 10 namespace blocks beginning at namespace block **430-a** (e.g., associated with index number 20) and including all sequentially ordered namespace blocks

until, and including, namespace block **435-a** (e.g., associated with index number 29). The fourth namespace can include two namespace blocks beginning at namespace block **440-a** (e.g., associated with index number 30) and include all sequentially ordered namespace blocks until, and including, namespace block **445-a** (e.g., associated with index number 31). The fifth namespace can include twelve namespace blocks beginning at namespace block **450** (e.g., associated with index number 32) and including all sequentially ordered namespace blocks until, and including, namespace block **455** (e.g., associated with index number 43). In this case, namespace block **460** can be the first available unmapped namespace block (e.g., associated with index number 44), and all namespace block sequentially ordered after namespace block **460**, until, and including namespace block **470-a** (e.g., associated with index number 59) can be namespace blocks available for namespace mapping. Although namespace blocks with index numbers between 44 and 59 are shown in namespace map **405-a** in increasing consecutive order, such an order is non-limiting and any ordering (e.g., a random order) of the namespace blocks may be considered without departing from the scope of the disclosure.

In the case of namespace map **405-a**, an indication **465** can be a variable (e.g., a private variable written to storage accessible by firmware) and can be associated with namespace block **460** (e.g., as described with reference to operation **235** of method **200-a**). In some cases, the indication **465** is a set of bits representative of the index of the namespace block with which the indication **465** is associated. The indication **465** indicates that namespace block **460** is the first available namespace block for namespace mapping. The indication **465** can indicate that all namespace blocks associated with index numbers greater than namespace block **460** are available for namespace mapping (e.g., associated with index numbers greater than 44).

As discussed regarding operation **240** of method **200-b**, the memory sub-system can receive a request to delete (e.g., unmap) a namespace. For example, the memory sub-system can receive a request to delete the fifth namespace (e.g., the fifth namespace including the twelve namespace blocks beginning at namespace block **450**, as associated with index number 32, and including all sequentially ordered namespace blocks until, and including, namespace block **455**, associated with index number 43).

As discussed regarding operation **245** of method **200-b**, the namespace can be unmapped. In this case, the namespace blocks previously associated with the now deleted namespace can be unmapped in the namespace map. For example, in response to the deletion of the fifth namespace, namespace blocks associated with index numbers 32 through 43 can be unmapped from the fifth namespace (e.g., disassociated from the fifth namespace), and be associated with unmapped (e.g., available) namespace blocks.

As discussed with reference to operation **250** of method **200-b**, after the unmapping of the namespace blocks from the deleted namespace, the indication **465** can be disassociated from the previously first-available namespace block. As illustrated in namespace map **405-a**, the indication **465** can have previously been written to namespace block **460** indicating that namespace block **460** was the first available namespace block. Upon deletion of the fifth namespace, the indication **465** can be disassociated with namespace block **460** as namespace block **460** is no longer be the first namespace block available for namespace mapping. In response to the unmapping of the fifth namespace (e.g.,

15

namespace blocks associated with index numbers 32 through 43 now being available for namespace mapping), and as discussed with reference to operation 255 of method 200-b, the available namespace blocks can be reordered. As discussed with reference to operation 260 of method 200-b, after reordering of the available namespace blocks, the indication 465 can be updated to be associated with a new first available namespace block. For example, after reordering of the available namespace blocks (e.g., as associated with index numbers 32 through 59), the new first available namespace block is namespace block 475 (e.g., associated with index number 32). In this case, an indication 465 can be associated with namespace block 475. The indication 465 can indicate that namespace block 475 is now the first available namespace block for namespace mapping.

FIG. 5 is an example machine of a computer system 500 in which examples of the present disclosure can operate. The computer system 500 can include a set of instructions, for causing the machine to perform any one or more of the techniques described herein. In some examples, the computer system 500 can correspond to a host system (e.g., the host system 105 described with reference to FIG. 1) that includes, is coupled with, or utilizes a memory sub-system (e.g., the memory sub-system 110 described with reference to FIG. 1) or can be used to perform the operations of a controller (e.g., to execute an operating system to perform operations corresponding to the namespace manager 150 described with reference to FIG. 1). In some examples, the machine can be connected (e.g., networked) with other machines in a LAN, an intranet, an extranet, and/or the Internet. The machine can operate in the capacity of a server or a client machine in client-server network environment, as a peer machine in a peer-to-peer (or distributed) network environment, or as a server or a client machine in a cloud computing infrastructure or environment.

The machine can be a personal computer (PC), a tablet PC, a set-top box (STB), a Personal Digital Assistant (PDA), a cellular telephone, a web appliance, a server, a network router, a switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while a single machine is illustrated, the term “machine” can also include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

The example computer system 500 can include a processing device 505, a main memory 510 (e.g., read-only memory (ROM), flash memory, DRAM such as SDRAM or RDRAM, etc.), a static memory 515 (e.g., flash memory, static random access memory (SRAM), etc.), and a data storage system 525, which communicate with each other via a bus 545.

Processing device 505 represents one or more general-purpose processing devices such as a microprocessor, a central processing unit, or the like. More particularly, the processing device can be a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets, or processors implementing a combination of instruction sets. Processing device 505 can also be one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. The processing device 505 is configured to execute instructions 535 for performing the

16

operations and steps discussed herein. The computer system 500 can further include a network interface device 520 to communicate over the network 540.

The data storage system 525 can include a machine-readable storage medium 530 (also known as a computer-readable medium) on which is stored one or more sets of instructions 535 or software embodying any one or more of the methodologies or functions described herein. The instructions 535 can also reside, completely or at least partially, within the main memory 510 and/or within the processing device 505 during execution thereof by the computer system 500, the main memory 510 and the processing device 505 also constituting machine-readable storage media. The machine-readable storage medium 530, data storage system 525, and/or main memory 510 can correspond to a memory sub-system.

In one example, the instructions 535 include instructions to implement functionality corresponding to a namespace manager 550 (e.g., the namespace manager 150 described with reference to FIG. 1). While the machine-readable storage medium 530 is shown as a single medium, the term “machine-readable storage medium” can include a single medium or multiple media that store the one or more sets of instructions. The term “machine-readable storage medium” can also include any medium that is capable of storing or encoding a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure. The term “machine-readable storage medium” can include, but not be limited to, solid-state memories, optical media, and magnetic media.

Some portions of the preceding detailed descriptions have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the ways used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. The present disclosure can refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage systems.

The present disclosure also relates to an apparatus for performing the operations herein. This apparatus can be specially constructed for the intended purposes, or it can include a general purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program can be stored in a computer readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic-optical disks, ROMs, RAMs,

17

EPROMs, EEPROMs, magnetic or optical cards, or any type of media suitable for storing electronic instructions, each coupled to a computer system bus.

The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general purpose systems can be used with programs in accordance with the teachings herein, or it can prove convenient to construct a more specialized apparatus to perform the method. The structure for a variety of these systems will appear as set forth in the description below. In addition, the present disclosure is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages can be used to implement the teachings of the disclosure as described herein.

The present disclosure can be provided as a computer program product, or software, that can include a machine-readable medium having stored thereon instructions, which can be used to program a computer system (or other electronic devices) to perform a process according to the present disclosure. A machine-readable medium includes any mechanism for storing information in a form readable by a machine (e.g., a computer). In some examples, a machine-readable (e.g., computer-readable) medium includes a machine (e.g., a computer) readable storage medium such as a ROM, RAM, magnetic disk storage media, optical storage media, flash memory components, etc.

In the foregoing specification, examples of the disclosure have been described with reference to specific example examples thereof. It will be evident that various modifications can be made thereto without departing from the broader scope of examples of the disclosure as set forth in the following claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

What is claimed is:

1. A method, comprising:
 - generating a namespace map having a set of namespace blocks, wherein an unmapped namespace block of the set of namespace blocks comprises an indication corresponding to a division of the namespace map where each namespace block following a sequential order before the unmapped namespace block associated with the indication is unavailable for mapping and each namespace block following the sequential order after the unmapped namespace block associated with the indication is available for mapping;
 - receiving a command to create a namespace using the namespace map; and
 - mapping, based at least in part on receiving the command, one or more available namespace blocks to the namespace, wherein the one or more available namespace blocks comprise at least the unmapped namespace block.
2. The method of claim 1, wherein a size of the unmapped namespace block is based at least in part on a quantity of logical block addresses associated with the unmapped namespace block.
3. The method of claim 1, wherein one or more logical block addresses associated with the unmapped namespace block are ordered within the unmapped namespace block.
4. The method of claim 1, wherein a size of the namespace map is based at least in part on a quantity of namespace blocks associated with the namespace map.
5. The method of claim 1, further comprising:
 - receiving a second command to erase the namespace; and

18

unmapping each namespace block associated with the namespace based at least in part on receiving the second command.

6. The method of claim 5, wherein unmapping each namespace block associated with the namespace comprises:
 - disassociating each namespace block from the namespace within the namespace map; and
 - associating each disassociated namespace block within the namespace map as namespace blocks that are available for namespace mapping based at least in part on an ordering of the namespace map.
7. The method of claim 6, further comprising:
 - updating a second indication of a first-available namespace block for mapping before mapping the unmapped namespace block;
 - disassociating the unmapped namespace block based at least in part on receiving the second command; and
 - updating the second indication with a second index of a second unmapped namespace block based at least in part on the unmapping of each namespace block.
8. The method of claim 7, further comprising:
 - storing the second indication at a memory sub-system, wherein the second indication is accessible by firmware of the memory sub-system.
9. The method of claim 5, further comprising:
 - reordering the one or more available namespace blocks according to a sequential order based at least in part on the unmapping of each namespace block.
10. The method of claim 1, further comprising:
 - associating a third indication with the unmapped namespace block, wherein the third indication is representative of a first-available namespace block for mapping;
 - disassociating the third indication from the unmapped namespace block after mapping the unmapped namespace block to the namespace; and
 - associating the third indication with a fourth unmapped namespace block based at least in part on the mapping of the one or more available namespace blocks to the namespace, wherein the fourth unmapped namespace block comprises a first available namespace block of available namespace blocks of the namespace map.
11. The method of claim 10, further comprising:
 - reordering the one or more available namespace blocks of the set of namespace blocks according to the sequential order based at least in part on the mapping of the one or more available namespace blocks.
12. The method of claim 1, wherein each namespace block of the set of namespace blocks is associated with respective groups of logical block addresses, and wherein an ordering of the namespace map is based at least in part on the respective groups of logical block addresses.
13. The method of claim 1, wherein the namespace map comprises a variable for indicating that the unmapped namespace block and each namespace block following the unmapped namespace block are available for mapping.
14. The method of claim 1, wherein the mapping of the one or more available namespace blocks to the namespace is performed upon boot-up of a memory sub-system.
15. A memory system, comprising:
 - a plurality of memory components; and
 - processing circuitry coupled with the plurality of memory components, wherein the processing circuitry is operable to cause the memory system to:
 - generate a namespace map having a set of namespace blocks, wherein an unmapped namespace block of the set of namespace blocks comprises an indication

19

corresponding to a division of the namespace map where each namespace block following a sequential order before the unmapped namespace block associated with the indication is unavailable for mapping and each namespace block following the sequential order after the unmapped namespace block associated with the indication is available for mapping; receive a command to create a namespace using the namespace map; and map, based at least in part on receiving the command, one or more available namespace blocks to the namespace, wherein the one or more available namespace blocks comprise at least the unmapped namespace block.

16. The memory system of claim **15**, wherein the processing circuitry is operable to cause the memory system to: receive a second command to erase the namespace; and unmap each namespace block associated with the namespace based at least in part on receiving the second command.

17. The memory system of claim **16**, wherein to unmap each namespace block associated with the namespace, the processing circuitry is operable to cause the memory system to:

disassociate each namespace block from the namespace within the namespace map; and associate each disassociated namespace block within the namespace map as namespace blocks that are available for namespace mapping based at least in part on an ordering of the namespace map.

18. A non-transitory computer-readable storage medium comprising instructions that, when executed by one or more processing devices, cause the one or more processing devices to:

generate a namespace map having a set of namespace blocks, wherein an unmapped namespace block of the

20

set of namespace blocks comprises an indication corresponding to a division of the namespace map where each namespace block following a sequential order before the unmapped namespace block associated with the indication is unavailable for mapping and each namespace block following the sequential order after the unmapped namespace block associated with the indication is available for mapping;

receive a command to create a namespace using the namespace map; and

map, based at least in part on receiving the command, one or more available namespace blocks to the namespace, wherein the one or more available namespace blocks comprise at least the unmapped namespace block.

19. The non-transitory computer-readable storage medium of claim **18**, wherein the instructions, when executed by the one or more processing devices, cause the one or more processing devices further to:

receive a second command to erase the namespace; and unmap each namespace block associated with the namespace based at least in part on receiving the second command.

20. The non-transitory computer-readable storage medium of claim **18**, wherein to unmap each namespace block associated with the namespace the instructions, when executed by the one or more processing devices, cause the one or more processing devices further to:

disassociate each namespace block from the namespace within the namespace map; and associate each disassociated namespace block within the namespace map as namespace blocks that are available for namespace mapping based at least in part on an ordering of the namespace map.

* * * * *