

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393521
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Bhoria; Naveen et al.

Methods and apparatus to facilitate write miss caching in cache system

Abstract

Methods, apparatus, systems and articles of manufacture to facilitate write miss caching in cache system are disclosed. An example apparatus includes a first cache storage; a second cache storage, wherein the second cache storage includes a first portion operable to store a first set of data evicted from the first cache storage and a second portion; a cache controller coupled to the first cache storage and the second cache storage and operable to: receive a write operation; determine that the write operation produces a miss in the first cache storage; and in response to the miss in the first cache storage, provide write miss information associated with the write operation to the second cache storage for storing in the second portion.

Inventors: Bhoria; Naveen (Plano, TX), Anderson; Timothy David (University Park, TX), Hippleheuser; Pete Michael (Murphy, TX)

Applicant: TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)

Family ID: 1000008766541

Assignee: TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)

Appl. No.: 18/337104

Filed: June 19, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230333991 A1	Oct. 19, 2023

Related U.S. Application Data

continuation parent-doc US 16882258 20200522 US 11693790 child-doc US 18337104
us-provisional-application US 62852494 20190524

Publication Classification

Int. Cl.: **G06F12/128** (20160101); **G06F9/30** (20180101); **G06F9/54** (20060101); **G06F11/10** (20060101); **G06F12/02** (20060101); **G06F12/0802** (20160101); **G06F12/0804** (20160101); **G06F12/0806** (20160101); **G06F12/0811** (20160101); **G06F12/0815** (20160101); **G06F12/0817** (20160101); **G06F12/0853** (20160101); **G06F12/0855** (20160101); **G06F12/0864** (20160101); **G06F12/0884** (20160101); **G06F12/0888** (20160101); **G06F12/0891** (20160101); **G06F12/0895** (20160101); **G06F12/0897** (20160101); **G06F12/1027** (20160101); **G06F12/12** (20160101); **G06F12/121** (20160101); **G06F12/126** (20160101); **G06F12/127** (20160101); **G06F13/16** (20060101); **G06F15/80** (20060101); **G11C5/06** (20060101); **G11C7/10** (20060101); **G11C7/22** (20060101); **G11C29/42** (20060101); **G11C29/44** (20060101)

U.S. Cl.:

CPC **G06F12/128** (20130101); **G06F9/3001** (20130101); **G06F9/30043** (20130101); **G06F9/30047** (20130101); **G06F9/546** (20130101); **G06F11/1064** (20130101); **G06F12/0215** (20130101); **G06F12/0238** (20130101); **G06F12/0292** (20130101); **G06F12/0802** (20130101); **G06F12/0804** (20130101); **G06F12/0806** (20130101); **G06F12/0811** (20130101); **G06F12/0815** (20130101); **G06F12/082** (20130101); **G06F12/0853** (20130101); **G06F12/0855** (20130101); **G06F12/0864** (20130101); **G06F12/0884** (20130101); **G06F12/0888** (20130101); **G06F12/0891** (20130101); **G06F12/0895** (20130101); **G06F12/0897** (20130101); **G06F12/1027** (20130101); **G06F12/12** (20130101); **G06F12/121** (20130101); **G06F12/126** (20130101); **G06F12/127** (20130101); **G06F13/1605** (20130101); **G06F13/1642** (20130101); **G06F13/1673** (20130101); **G06F13/1689** (20130101); **G06F15/8069** (20130101); **G11C5/066** (20130101); **G11C7/10** (20130101); **G11C7/1015** (20130101); **G11C7/106** (20130101); **G11C7/1075** (20130101); **G11C7/1078** (20130101); **G11C7/1087** (20130101); **G11C7/222** (20130101); **G11C29/42** (20130101); **G11C29/44** (20130101); **G06F2212/1016** (20130101); **G06F2212/1021** (20130101); **G06F2212/1024** (20130101); **G06F2212/1041** (20130101); **G06F2212/1044** (20130101); **G06F2212/301** (20130101); **G06F2212/454** (20130101); **G06F2212/603** (20130101); **G06F2212/6032** (20130401); **G06F2212/6042** (20130101); **G06F2212/608** (20130101); **G06F2212/62** (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5003459	12/1990	Ramanujan et al.	N/A	N/A
5577227	12/1995	Finnell et al.	N/A	N/A
5644753	12/1996	Ebrahim	N/A	N/A
5687338	12/1996	Boggs et al.	N/A	N/A
6078992	12/1999	Hum	N/A	N/A
6170040	12/2000	Lee et al.	N/A	N/A
6226713	12/2000	Mehrotra	N/A	N/A

6513104	12/2002	Gaskins	N/A	N/A
6775750	12/2003	Krueger	N/A	N/A
8181005	12/2011	Zuraski, Jr. et al.	N/A	N/A
8341353	12/2011	Venkumahanti et al.	N/A	N/A
9170955	12/2014	Forsyth et al.	N/A	N/A
2004/0193808	12/2003	Spencer	N/A	N/A
2006/0143396	12/2005	Cabot	N/A	N/A
2009/0198867	12/2008	Guthrie et al.	N/A	N/A
2010/0023695	12/2009	Guthrie et al.	N/A	N/A
2011/0082981	12/2010	Hoogerbrugge	N/A	N/A
2011/0173391	12/2010	Venkumahanti et al.	N/A	N/A
2012/0042126	12/2011	Krick et al.	N/A	N/A
2012/0221774	12/2011	Atkisson	N/A	N/A
2012/0221793	12/2011	Tran	N/A	N/A
2013/0191601	12/2012	Peterson et al.	N/A	N/A
2014/0189245	12/2013	Rupley et al.	N/A	N/A
2014/0281248	12/2013	Alameldeen	N/A	N/A
2015/0006820	12/2014	Bhoria et al.	N/A	N/A
2016/0196210	12/2015	Noguchi et al.	N/A	N/A
2017/0177500	12/2016	Shanbhogue et al.	N/A	N/A
2017/0293561	12/2016	Dwiel et al.	N/A	N/A
2018/0024931	12/2017	Moyer	N/A	N/A
2018/0089091	12/2017	Akenine-Moller et al.	N/A	N/A
2019/0018799	12/2018	Appu et al.	N/A	N/A
2019/0384718	12/2018	Hansson et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1979451	12/2006	CN	N/A
1991792	12/2009	CN	N/A
102165424	12/2010	CN	N/A
101558391	12/2012	CN	N/A
108604210	12/2021	CN	N/A
108334458	12/2022	CN	N/A

OTHER PUBLICATIONS

International Search Report for PCT/US2020/034560 mailed Aug. 20, 2020. cited by applicant
International Search Report for PCT/US2020/034557 mailed Sep. 10, 2020. cited by applicant
Extended European Search Report for 20815004.5; dated Jun. 21, 2022. cited by applicant
Wang, et al. "Adaptive Placement and Migration Policy for an STT_RAM-Based Hybrid Cache";
2014 IEEE 20th International Symposium on High Performance Computer Architecture; XP
032606782; Feb. 15, 2014; pp. 13-24. cited by applicant
Office Action in corresponding China App. No. 202080038198.6, issued Jun. 6, 2025. cited by
applicant

Primary Examiner: Choi; Charles J

Background/Summary

RELATED APPLICATIONS (1) This patent is a continuation of U.S. patent application Ser. No. 16/882,258, filed May 22, 2020, which claims the benefit of and priority to U.S. Provisional Patent Application No. 62/852,494, which was filed on May 24, 2019, each of which is incorporated by reference herein in its entirety.

FIELD OF THE DISCLOSURE

(1) This disclosure relates generally to computer architecture, and, more particularly, to methods and apparatus to implement a data cache.

BACKGROUND

(2) Computing systems include one or more processing cores to execute instructions by accessing data stored in memory. However, the amount of time it takes for the processing core to access data from the memory can be significant. Accordingly, most computing systems include a cache which stores an amount of data from the memory (e.g., typically smaller than the total amount of data in the memory) that has a high probability of being accessed by the processing core in the future. Accordingly, when the processing core calls for data, the cache can provide the data to the processing core faster than the processing core retrieving the data from the memory, thereby increasing the speed and efficiency of the computing system.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is an example computing system described in conjunction with examples disclosed herein.
- (2) FIG. 2 is a block diagram of an example level one cache of the example computing system of FIG. 1.
- (3) FIGS. 3A-3D illustrate an example circuit implementation of the level one cache of the example computing system of FIG. 1.
- (4) FIG. 4A is an example circuit implementation of the main cache store queue of FIGS. 3A-3D.
- (5) FIG. 4B is an example circuit implementation of example data forwarding logic implemented in connection with the main cache store queue of FIGS. 3A-4A.
- (6) FIG. 4C is an additional example circuit implementation of the main cache store queue and/or of a victim cache store queue.
- (7) FIG. 4D is an example hardware implementation of a merging circuit of FIG. 4A.
- (8) FIG. 5 illustrates an example victim cache-based section of the level one cache of FIGS. 3A-3D.
- (9) FIG. 6 illustrates an example first table and an example second table for implementing an example replacement policy component of FIGS. 3A-3D and 5.
- (10) FIG. 7A illustrates an example encapsulated data cache system including a victim cache store queue and victim storage.
- (11) FIG. 7B illustrates another example encapsulated data cache system including a main cache store queue and main storage.
- (12) FIG. 7C illustrates an example victim cache bank structure.
- (13) FIG. 7D illustrates an example main cache bank structure.
- (14) FIG. 7E illustrates an example unified cache bank structure.
- (15) FIG. 8A illustrates a schematic illustration of an example implementation of victim cache tag random access memory.
- (16) FIG. 8B illustrates another schematic illustration of the example implementation of the victim cache tag random access memory of FIG. 8A.
- (17) FIGS. 9A and 9B (collectively FIG. 9) illustrate example instruction servicing workflows.

(18) FIGS. 10A-1 and 10A-2 (collectively FIG. 10A) illustrate an example data cache system.

(19) FIGS. 10B-1 and 10B-2 (collectively FIG. 10B) illustrate a second example data cache system.

(20) FIG. 11A is an example circuit implementation of the victim cache store queue of FIGS. 3A-3D.

(21) FIGS. 11B-1 and 11B-2 (collectively FIG. 11B) illustrate an example circuit implementation of an example unified cache store queue.

(22) FIG. 12-33 are flowcharts representative of machine readable instructions which may be executed to implement the level one cache of FIGS. 1-3D.

(23) FIG. 34 is a block diagram of an example processing platform structured to execute the instructions of FIGS. 3A-3D to implement the level one cache of FIGS. 1-3D.

DETAILED DESCRIPTION

(24) The figures are not to scale. Instead, the thickness of the layers or regions may be enlarged in the drawings. In general, the same reference numbers will be used throughout the drawing(s) and accompanying written description to refer to the same or like parts. As used in this patent, stating that any part (e.g., a layer, film, area, region, or plate) is in any way on (e.g., positioned on, located on, disposed on, or formed on, etc.) another part, indicates that the referenced part is either in contact with the other part, or that the referenced part is above the other part with one or more intermediate part(s) located therebetween. Connection references (e.g., attached, coupled, connected, and joined) are to be construed broadly and may include intermediate members between a collection of elements and relative movement between elements unless otherwise indicated. As such, connection references do not necessarily infer that two elements are directly connected and in fixed relation to each other. Stating that any part is in “contact” with another part means that there is no intermediate part between the two parts. Although the figures show layers and regions with clean lines and boundaries, some or all of these lines and/or boundaries may be idealized. In reality, the boundaries and/or lines may be unobservable, blended, and/or irregular.

(25) Various forms of the term “couple” are used throughout the specification. These terms may cover connections, communications, or signal paths that enable a functional relationship consistent with the description of the present disclosure. For example, if device A generates a signal to control device B to perform an action, in a first example device, A is coupled to device B by direct connection, or in a second example device, A is coupled to device B through intervening component C if intervening component C does not alter the functional relationship between device A and device B such that device B is controlled by device A via the control signal generated by device A.

(26) Consistent with the present disclosure, the term “configured to” describes the structural and functional characteristics of one or more tangible non-transitory components. For example, a device that is “configured to” perform a function mean that the device has a particular configuration that is designed or dedicated for performing a certain function. A device is “configured to” perform a certain function if such a device includes tangible non-transitory components that can be enabled, activated, or powered to perform that certain function. While the term “configured to” may encompass being configurable, this term is not limited to such a narrow definition. Thus, when used for describing a device, the term “configured to” does not require the described device to be configurable at any given point of time.

(27) Moreover, the term “example” is used herein to mean serving as an instance, illustration, etc., and not necessarily as advantageous. Also, although the disclosure has been shown and described with respect to one or more implementations, equivalent alterations and modifications will be apparent upon a reading and understanding of this specification and the annexed drawings. All such modifications and alterations are fully supported by the disclosure and is limited only by the scope of the following claims. In particular regard to the various functions performed by the above described components (e.g., elements, resources, etc.), the terms used to describe such components

are intended to correspond, unless otherwise indicated, to any component which performs the specified function of the described component (e.g., that is functionally equivalent), even though not structurally equivalent to the disclosed structure. In addition, while a particular feature of the disclosure may have been disclosed with respect to only one of several implementations, such feature may be combined with one or more other features of the other implementations as may be desired and advantageous for any given or particular application.

(28) While this specification contains many specifics, these should not be construed as limitations on the scope of what may be claimed, but rather as descriptions of features that may be specific to particular embodiments. Certain features that are described in this specification in the context of separate embodiments can also be implemented in combination in a single embodiment.

Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable subcombination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a subcombination or variation of a subcombination.

(29) Similarly, while operations are depicted in the drawings in an example particular order, this does not require that such operations be performed in the example particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results unless such order is recited in one or more claims. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above does not require such separation in all embodiments.

(30) Descriptors “first,” “second,” “third,” etc. are used herein when identifying multiple elements or components which may be referred to separately. Unless otherwise specified or understood based on their context of use, such descriptors do not impute any meaning of priority, physical order or arrangement in a list, or ordering in time but are merely used as labels for referring to multiple elements or components separately for ease of understanding the disclosed examples. In some examples, the descriptor “first” may be used to refer to an element in the detailed description, while the same element may be referred to in a claim with a different descriptor such as “second” or “third.” In such instances, it should be understood that such descriptors are used merely for ease of referencing multiple elements or components.

(31) A central processing unit (CPU) is electronic circuitry that executes instructions making up a program or workload. A CPU may include one or more processor cores to execute the instructions by accessing data from main memory (e.g., extended memory). Because it takes time to access data from main memory, the one or more processor cores may obtain data from the memory and store it locally in local memory (e.g., a data cache). The local memory is smaller and faster than the main memory. In this manner, the processor cores can use and/or manipulate the data locally rather than interfacing with the main memory. Because the latency (e.g., the amount of time needed to read and/or write data to/from memory) for reading and/or writing data is smaller when accessing the cache, storing data from the extended memory in the cache increases the speed and efficiency of a computing system.

(32) Some local memory devices (e.g., caches) include one or more victim caches. A victim cache is an additional storage included in or connected to a cache. Victim caches improve (e.g., reduce) the cache miss rate, and particularly reduce conflict misses, by storing data recently evicted from the corresponding cache. The addition of a victim cache can have a similar impact on cache performance. The benefit is most evident in cases when a victim cache is added to a direct mapped cache, because a direct mapped cache has a relatively high rate of conflict misses.

(33) Some examples disclosed herein include an improved data cache and victim cache architecture that reduces overall cache latency. In some alternative cache designs, a victim cache typically resides at the end of a cache controller. When a line is evicted out of the cache controller (e.g.,

removed to make room for additional data), instead of removing the data line and sending the evicted data line all the way to the next level cache (e.g., L2) to be stored, the victim cache holds the evicted data line (e.g., victim data) until there is a reason to evict it out due to capacity or other reasons. In such a system, when a subsequent miss in the corresponding cache (e.g., L1 cache) occurs, the victim cache can service the miss and return the line to the corresponding cache, and thereby to a requesting processor core, so that the latency of read-miss is reduced. However, because some example systems include a victim cache in series with a main cache, the victim cache has to wait for the main cache to determine that the main cache does not have a data corresponding to a memory operation (also referred to herein as a memory instruction, memory request, and/or memory transaction) from a CPU (e.g., a read operation, instruction, request, and/or transaction, a write operation, instruction, request, and/or transaction, a read-modify-write operation, instruction, request, and/or transaction, an atomic operation, instruction, and/or transaction, a histogram operation, instruction, request, and/or transaction, etc.) before a victim cache can determine whether the memory address is stored in the victim cache. Examples disclosed herein include a cache system that decreases latency (e.g., increases speed), reduces the number of cycles to perform write operations, atomic operations (e.g., manipulating data stored in a location to increment, decrement, swap with other data, etc.), and/or histogram operations (e.g., determine a total number of each value in a cache line), and/or increases efficiency of a computing system that includes a cache.

(34) FIG. 1 illustrates an example computing system **100**. The example computing system **100** includes an example CPU **102**, example processing cores **104a-104n**, an example extended memory **106**, and an example data cache **108**. The example data cache **108** includes an example level one (L1) cache **110**, an example level two (L2) cache **112**, and an example level three (L3) cache **114**. Although the example computing system **100** of FIG. 1 includes N processing cores and three levels of cache. The example computing system **100** may include any number of processing cores and/or levels of cache. Additionally, one or more of the example components of the computing system **100** may be implemented on the same die and/or different dies in the same chip and/or different chips.

(35) The example CPU **102** of FIG. 1 includes the example processing cores **104a-104n**. The processing cores **104a-104n** may be incorporated in the same die or in separate dies (e.g., connected via one or more interfaces). The CPU **102** connects to the example data cache **108** via one or more interfaces. For example, the example core **104a** of the example CPU **102** may be connected to the data cache **108** via a scalar interface (e.g., a 64-bit wide bidirectional and/or unidirectional interface) and/or a vector interface (e.g., a 512-bit wide interface). The use of the scalar interface and the vector interface may be based on interface utilization, data size, and/or other considerations. For example, a scalar read operation may be transmitted via the scalar data while the data corresponding to the scalar read may be transmitted over the vector interface. Additionally, either the scalar interface and/or the vector interface may be used if the other interface is busy. In some examples, the CPU **102** may be connected to the data cache **108** using a different number and/or different types of interfaces.

(36) The processing cores **104a-104n** of FIG. 1 access data from the example extended memory **106** and/or the example cache **108** to execute instructions stored in the example extended memory **106** and/or the example data cache **108**. The example cores **104a-104n** can execute different processes at the same time to speed up the execution of tasks (e.g., instructions) of the computing system **100**. For example, the cores **104a-104n** may execute instructions from a workload of an application by writing and/or reading data to/from the extended memory **106** and/or the data cache **108**. Because the data cache **108** includes a copy of some of the data stored in the example extended memory **106**, when one of the example cores **104a** needs to access (e.g., read and/or write) data from the extended memory **106**, the example core **104a** transmits the read and/or write instructions to the example data cache **108**. As further described below, if the data cache **108**

includes the data corresponding to the instructions from the core **104a** (e.g., corresponding to a cache hit), the data cache **108** fulfills the request and/or instructions from the processing core **104a**. If the data cache **108** does not include the data corresponding to the instructions from the cores **104a** (e.g., corresponding to a cache miss), the data cache **108** interfaces with the example extended memory **106** to perform the transaction from the core **104a**.

(37) The example extended memory **106** of FIG. **1** is connected to the example data cache **108** via an interface (e.g., a 1024-bit wide extended memory interface). However, a different number and/or a different type of interface may be additionally or alternatively used. The example extended memory **106** stores all the data that can be accessed for the computing system **100**. The example extended memory **106** can be incorporate into the computing system **100** and/or may be memory external to the computing system **100** (e.g., off-chip memory). In some examples, extended memory **108** interacts with a controller of the highest cache (e.g., the example L3 cache **114**), and the highest cache (e.g., L3 cache **114**) interacts with the next highest cache (e.g., the example L2 cache **112**), and so on. In such examples, the CPU **102** transmits memory operations to the example L1 cache **110** and if the memory operation cannot be served by the L1 cache **110**, the L1 cache **110** transmits the memory operation to L2 cache **112**, and so on. In the example of FIG. **1**, if the L3 cache **114** (e.g., the highest level cache) cannot perform a read or write transaction (e.g., because the memory address is not located in the local storage of the L3 cache **114**), the L3 cache **114** interacts with the extended memory **106** to read or write the corresponding data to the memory address. The extended memory **106** may be on chip or off chip memory (e.g., DDR) and the interface to the extended memory may be 2.sup.N bits, where N depends on the type of extended memory used. In some examples, there can be a prefetcher and/or preload mechanism in any of the data caches which can pull data from the example extended memory **106** prior to execution of a problem to be stored locally at the cache before the CPU **102** executes any instructions, the memory **106** provides copies of the data stored in the memory to the example data cache **108**. The data cache **108** may request additional information and/or instruct the extended memory **106** to adjust the stored data in the extended memory **106** periodically, aperiodically, and/or based on a trigger, based on instructions from the CPU **102**.

(38) The example data cache **108** of FIG. **1** stores blocks of data (e.g., a cached subset of the data stored in the extended memory **106**) from the example extended memory **106** to reduce the time needed for the example CPU **102** to access the cached subset, thereby improving system performance. For best performance, attempts are made so that the data in the data cache **108** corresponds to the data most likely to be used by the CPU **102**. The data cache **108** provides access to the cached data when called upon by the CPU **102** during a cache hit (e.g., when the requested data is stored in the data cache **108**). If the CPU **102** requests data that is not included in the data cache **108** (e.g., a cache miss), the data cache **108** retrieves the corresponding data from the extended memory **106**. For example, if the CPU **102** transmits a read operation for corresponding to a particular address of the extended memory **106**, the data cache **108** determines whether the data corresponding to the particular address of the extended memory **106** is cached in the data cache **108**. If the data cache **108** determines that the data is cached, the data cache **108** provides the data to the CPU **102** to be read. If the data cache **108** determines that the data is not cached, the data cache **108** obtains the data from the extended memory **106** and returns the data to the CPU **102**. Additionally, the example data cache **108** may update the cached data based on the cache misses (e.g., to mitigate against future cache misses). For write transactions from the CPU **102**, the example data cache **108** may temporarily store the corresponding data before providing it to update older data stored in the extended memory **106**.

(39) The example data cache **108** includes the example L1 cache **110**, the example L2 cache **112**, and the example L3 cache **114**. The levels of the cache may be based on speed and/or size. For example, the example L1 cache **110** may be the fastest cache and smallest, followed by L2 **112** (e.g., slower than L1 **110** but larger) and L3 **114** (e.g., slower than L2 **112** but larger). Typically, the

most commonly used data is stored in the L1 **110** followed by the L2 **112** and the L3 **114**.

Accordingly, when the CPU **102** calls for an address to be written to or read from, the instruction from the CPU **102** is first sent to the L1 cache **110** and, if the corresponding data is not stored in the L1 cache **110**, then the instruction is sent to the L2 cache **112**. If the corresponding data is not stored in the L2 cache **112**, the instruction is sent to the L3 cache **114**. If the corresponding data is not stored in the L3 cache **114**, the example data cache **108** accesses the data from the extended memory **106**.

(40) As further described below, the example L1 cache **110** of FIG. **1** includes a main cache and victim cache to execute instructions from the CPU **102** in parallel. A victim cache stores victim data (e.g., data that was evicted from the main cache to make room for new data corresponding to other address locations of the extended memory **106**). In this manner, when there is a cache miss at the L1 level, the main L1 cache evicts a first set of data (e.g., the victim data) to make room for a second set of data that corresponds to the transaction that caused the cache miss. The victim cache then stores the first set of data that was removed from the main cache to make room for the second set of data in the main cache.

(41) FIG. **2** is an example block diagram of the L1 cache **110** of FIG. **1**. The example L1 cache **110** includes an example CPU interface **202**, example tag random access memory (RAM) accesses **204**, **206**, example tag RAM **208**, **210**, an example main cache store queue **212**, an example main storage **214** (also referred to as main cache storage or main cache), an example victim cache store queue **216**, an example victim storage **218** (also referred to as victim cache storage or victim cache), an example cache controller **220**, an example main cache controller **222**, an example victim cache controller **224**, and an example L2 interface **228**.

(42) The example CPU interface **202** connects the CPU **102** (e.g., the core(s) **104a-104n** of the CPU **102**) to the components of the L1 cache **110**. The CPU interface **202** is connected to the tag RAM accesses **204**, **206** and the cache controller **220**. The example CPU interface **202** receives instructions from the example core(s) **104a-104n** of the example CPU **102**. The instructions may include, read instructions, write instructions, read-modify-write instructions, atomic instructions, etc. When the CPU interface **202** obtains instructions corresponding to particular data stored at a particular address, the CPU interface **202** interfaces with the cache controller **220** and the main tag RAM access **204** to determine whether the corresponding data is stored in the main storage **214** and/or the victim storage **218** to perform the transaction. Additionally, for some types of transactions (e.g., read transactions) the example CPU interface **202** returns corresponding data to the example CPU **102**. Furthermore, when the CPU interface **202** obtains instructions corresponding to data at a particular address, the main components (e.g., the example main tag RAM access **204**, the example tag RAM **208**, the example main cache store queue **212**, the example main storage **214**, and the example main cache controller **222**) and the victim components (e.g., the example tag RAM access **206**, the example tag RAM **210**, the example victim cache store queue **216**, the example victim storage **218**, and the example victim cache controller **224**) operate in parallel (e.g., at the same time) to determine whether data from the corresponding address is stored in the main storage **214** and/or the victim storage **218**.

(43) The example main tag RAM access **204** of FIG. **1** is coupled to the tag RAM **208** and the cache controller **220**. The victim tag RAM access **206** is coupled to the tag RAM **210** and the cache controller **220**. The main tag RAM access **204** accesses the tag RAM **208** to determine whether the data from a memory address corresponding to the instructions from the CPU **102** is present in the main storage **214**. The example victim tag RAM access **206** accesses the tag RAM **210** to determine whether the data from a memory address corresponding to the instructions from the CPU **102** is present in the victim storage **218** in parallel with the main tag RAM access **204**. In some examples, the main tag RAM access **204** is implemented in the tag RAM **208** and the victim tag RAM access **206** is implemented in the tag RAM **210**. When the main tag RAM access **204** and/or the victim tag RAM access **206** determines address(es) corresponding to the instructions from the

CPU **102** is/are present in the respective tag RAM **208,210**, the main tag RAM access **204** and/or the victim tag RAM access **206** transmits the results (e.g., the determination and/or any corresponding data) to the example cache controller **220**.

(44) In some examples, the main storage is directly mapped. Accordingly, in such examples, a particular CPU memory address can only be stored at a particular location of the main storage **214**. Thus, the example tag RAM **208** can potentially have a fixed memory address for the CPU instruction. In a directly mapped cache, a given address is stored at a particular location of the tag RAM **208**.

(45) The example tag RAM **208** of FIG. 2 is coupled to the example cache controller **220** and the example main storage **214**. The example tag RAM **208** stores a table that records the entries in the example main storage **214** that correspond to memory addresses in the extended memory **106**. In this manner, the example main tag RAM access **204** can review the table to determine if data corresponding to instructions from the CPU **102** is available in the main storage **214**. The example tag RAM **210** is coupled to the example cache controller **220** and the example victim storage **218**. The example tag RAM **210** stores a table that records the entries in the example victim storage **218**. In this manner, the example victim tag RAM access **206** can review the table to determine if data corresponding to instructions from the CPU **102** is available in the victim storage **218**. When the data of the main storage **214** is updated, the example tag RAM **208** may update the table to reflect any changes to the entries. For example, if the main storage **214** removes a first set of data corresponding to a first memory address in the extended memory **106** and replaces it with a second set of data corresponding to a second memory address in the extended memory **106**, the example tag RAM **208** updates the table to reflect that the first set of data of the first memory address is no longer stored in the main storage **214** and the second set of data of the second memory address is now stored in the main storage **214**. Likewise, when data from a first address location of the victim storage **218** is changed to data from a second address location, the example tag RAM **210** may update the table to reflect any changes to the entries.

(46) The example victim-side tag RAM **210** may be a content addressable memory (CAM). In some examples, the victim storage **218** is fully-associative (e.g., any location of the victim storage **218** can be used to store data from any CPU address). Thus, when the example CPU **102** provides a memory address to the example L1 cache **110**, the example victim tag RAM **210** compares the provided memory address to all the entries of the tag RAM **210**. If there is a match between the provided address and the entries stored in the tag RAM **210**, then the address of the corresponding location in the victim storage **218** is output by the tag RAM **210**. The address is used to obtain the data from the victim storage **218** that corresponds to the CPU instruction.

(47) The example main cache store queue **212** of FIG. 2 is coupled to the example main storage **214** and the example cache controller **220**. The example main cache store queue **212** is used when the CPU **104** issues a store operation (e.g., write operation, atomic compare and swap, atomic operation, etc.). The example main cache store queue **212** can implement read-modify-write functionality. Read-modify-write functionality includes storing data in the local storages where there is a need to first read an older version of the data being written that is already present in the main storage **214** for reasons including updating error correction code data. In this manner, while the read operation is being performed to get a copy of the older data from the main storage **214**, the new bit(s) included in the write portion of the read-modify-write transaction from CPU **102** is/are buffered in the main cache store queue **212**. When the older data is available then only the new bit(s) which the CPU **102** is/are overwriting are updated in the store queue buffer and the updated (e.g., merged) data is written back to the main storage **214**. For example, the new bit(s) which the CPU **102** is/are overwriting from the write portion of the read-modify-write transaction (e.g., from the CPU **102**) including the corresponding memory address are buffered in the main cache store queue **212** until the old data of the corresponding write (e.g., store) instruction is read from the main storage **214**. When the new bit(s) included in the write portion of the read-modify-write

transaction are merged with the old data from the main storage **214**, the updated (e.g., merged) data is written back into the main storage **214**. The extra pipelining for the store instructions provided by the main cache store queue **212** enables the instructions from the CPU **104** to continue execution when prior store instructions are waiting for load data from the main storage **214** to be available. In some examples, a main storage **214** includes SRAM, as further described below in conjunction with Sections 5 and 16. The example main cache store queue **212** is further described below in conjunction with FIGS. **3A**, **3B**, and **4**.

(48) The example victim cache store queue **216** of FIG. **2** is coupled to the example victim storage **218** and the example cache controller **220**. The example victim cache store queue **216** buffers data that is to be stored in the example victim storage **218**. For example, victim values from store instructions (e.g., from the cache controller **220**) including the corresponding memory address and/or store data are buffered in the victim cache store queue **216** until a prior or old value of the corresponding store instruction is read from the victim storage **218**. When stored bytes are merged with the prior/old data from the victim storage **218**, the value is written back into the victim storage **218**. While the victim cache store queue **216** may process read, modify, and/or write operations from the cache controller **220** that were transmitted in response to a retirement point met (e.g., when one or more cache lines is removed from the L1 cache **110** to the L2 cache **112**), in other examples disclosed herein, the victim cache store queue **216** may process read, modify, and/or write operations from the cache controller **220** that were transmitted directly from the CPU **102**. The example victim cache store queue **216** is further described below.

(49) The example main storage **214** of FIG. **2** is coupled to the example tag RAM **208**, the example main cache store queue **212**, and the example cache controller **220**. The example main storage **214** stores data (e.g., entries) that correspond to data in memory address locations in the extended memory **106**. The main storage **214** stores that data that has a high probability of being called upon by the example CPU **102**. The stored data may be updated when, for example, the CPU **102** attempts to access (e.g., read, write, etc.) data that is not cached in the example main storage **214**. The example main storage **214** may be or include static RAM (SRAM) and/or any other type of storage which provides a single cycle access to stored data.

(50) The example victim storage **218** of FIG. **2** is coupled to the example the example tag RAM **210**, the example victim cache store queue **216** and the example cache controller **220**. The example victim storage **218** stores data (e.g., entries) corresponding to memory address locations that have been removed from the example main storage **214** to make room for data from other address locations that the CPU **102** more recently accessed or is more likely to access. The example victim storage **218** may be or include a register file, static RAM (SRAM), and/or any other type of storage.

(51) In operation, initially, the example main storage **214** stores data from different addresses in the example extended memory **106** based on likelihood of use. If the example CPU **102** sends an instruction corresponding to a memory address that is stored in the main storage **214**, the example cache controller **220** controls the components to perform the transaction. For example, the cache controller **220** may transmit an instruction to the main cache store queue **212** for subsequent processing and storage in the main storage **214**. If the example CPU **102** sends an instruction corresponding to a memory address that is not stored in the main storage (e.g., stored in the victim storage **218**), the example main storage **214** may store the data at the corresponding memory address after the data is retrieved from another storage. To make room for the data, the main storage **214** removes some data corresponding to one or more memory address locations. The removed data and corresponding memory address locations are called victims or evicted lines. The selection of which data is to be a victim is based on various parameters according to a replacement policy. After being removed from the main storage **214**, the victim is stored in the example victim storage **218**. If the victim storage **218** is full, the victim storage **218** removes data (e.g., a second victim) to make room to store the victim from the main storage **214**. After the second victim is

removed from the victim storage **218**, the second victim is transmitted to L2 cache **112** for storage. (52) Although the example L1 cache **110** of FIG. **1** includes one main cache store queue **212** with corresponding main storage **214** and one example victim cache store queue **216** with corresponding victim storage **218**, the example main cache store queue **212** and/or the example victim cache store queue **216** may include multiple store queues corresponding to multiple storages. For example, the main storage **214** and/or the victim storage **218** may be divided into a plurality of independently addressable banks (e.g., into any number of storage devices with any number of line widths, rows, etc.), where individual banks may have its own corresponding store queue. For example, a main storage **214** made up of 256 rows with each row having a line width of 1024 bits may be divided into 16 main storages (e.g., banks) with 64 bits per row in a particular bank, where each bank may have their own main store queue. In such an example, if vector data for a read and/or write request that hits the L1 cache **110** is 512 bits wide, the cache controller **220** may process the request as 8 parallel writes and/or reads to the 8 banks. In such a multi-bank set up, the read and/or write operations may be sent to the banks in parallel and the banks arbitrate their own process in response to the read and/or write operations. By operating independently, operation of the multi-bank approach is more efficient than a single main storage (e.g., because the entire cache line is not locked up when a request is received and only a portion of the cache line allocated to the bank that received such a request would be locked), at the expense of complication and/or parts.

(53) The example cache controller **220** of FIG. **2** is coupled to the components of the L1 to control how data is read and/or written in the example storages **214**, **216**, and/or how data is updated in the example storages **214**, **218**. For example, when a read request, a write request, an atomic request, a read-modify-write request, etc. is received at the example CPU interface **202**, the cache controller **220** obtains the request and instructs the other components accordingly. For example, during a read request for data at a particular location of the extended memory **106**, the example cache controller **220** instructs the main tag RAM access **204** to access the tag RAM **208** to determine if the main storage **214** is storing the data corresponding to the location of the extended memory **106** from the read request. If the main tag RAM access **204** determines that the entry is located in the main storage **214** (e.g., a cache hit), the cache controller **220** determines the location of the particular entry based on the data in the tag RAM **208** and interfaces with the main cache store queue **212** to read the value from the example main storage **214** and returns the value to the CPU **102** via the CPU interface **202**. The example cache controller **220** includes the main cache controller **222** to control the main cache components (e.g., the example main tag RAM access **204**, the example tag RAM **208**, the example main cache store queue **212**, and the example main storage **214**) and the example victim cache controller **224** to control the victim cache components (e.g., the example victim tag RAM access **206**, the example tag RAM **210**, the example victim cache store queue **216**, and the example victim storage **218**) in parallel with the main cache controllers **222** control of the main cache components. In some examples, the cache controllers **222**, **224** may be separate controllers and/or combined to operate in a single controller.

(54) The example cache controller **220** of FIG. **2** interfaces with the example L2 interface **228** to obtain data to be stored in the example main storage **214** (e.g., initially, after a cache miss, etc.). Additionally, the example cache controller **220** may transmit updates to the data of the main storage **214** and/or the victim storage **218** to the L2 cache **112** when new data is written (e.g., when old data is written over), so that read instructions from the CPU **102** are propagated to the L2 cache **112** via the L2 interface **228**. The example cache controller **220** interfaces with the example L2 interface **228** to transmit and/or receive data from the L2 cache **112** and/or the L3 cache **114** (e.g., directly or via the L2 cache **112**). For example, when the main storage **214** and the victim storage **218** do not include a memory address location for data corresponding to instructions from the CPU **102** (e.g., corresponding to a cache miss), the example cache controller **220** transmits the instruction to the example L2 cache **112** via the L2 interface **228**. The cache controller **220** may receive the corresponding data from the L2 cache **112** via the L2 interface **228** if the data is stored

in the L2 cache **112** or the L3 cache **114**. In such an example, the cache controller **220** may store the corresponding information from the L2 cache **112** in the main storage **214**. In some examples, when the victim storage **218** has to remove data for a particular location (e.g., an old victim) to make room for a new victim from the main storage **214**, the example cache controller **220** may transmit the old victim to the L2 cache **112** via the L2 interface **228** to be stored in the L2 cache.

(55) FIGS. 3A-3D illustrate an example circuit implementation of the L1 cache **110** of the example computing system **100** of FIG. 1. The example implementation of FIGS. 3A-3D includes the example CPU interface **202**, the example tag RAMs **208**, **210**, the example main cache store queue **212**, the example main storage **214**, the example victim cache store queue **216**, the example victim storage **218**, and the example cache controller **220** of FIG. 2. The example implementation of FIGS. 3A-3D further includes an example modify, exclusive, shared, invalid (MESI) RAM **300**, example address processing components **302a-c**, example bank processing logic **303**, example hit/miss comparison logic **304**, **306**, an example replacement policy component **308**, an example flush engine **309**, an example error correcting code (ECC) logic **310**, **312**, example data multiplexer (MUX) circuits **314**, **316**, an example MUXs **318**, **320**, and example latches **322**.

(56) In the example implementation of FIGS. 3A-3D, the example CPU interface **202** includes two interfaces (e.g., one scalar and one vector interface, both interfaces having two parts, one for input data from the CPU **102** and one for output data to the CPU **102**). The input CPU interface **202** of FIGS. 3A-3D includes an elastic buffer to buffer incoming data from the CPU **102**, a multiplexer to select between the buffered data from an elastic buffer in case there are pending CPU instructions in the elastic buffer and instructions coming directing from the CPU **102** in case the elastic buffer queue is empty, and breaks the incoming instructions into the corresponding address, operation (e.g., read, write, etc.) and write data (e.g., if the instructions correspond to a write operation). The output CPU interface **202** of FIGS. 3A-3D transmits data back to the CPU **102**.

(57) The example main cache store queue **212** of FIGS. 3A-3D includes blocks that correspond to operations of the main cache store queue **212**. For example, the main cache store queue **212** includes blocks to implement a read-modify-write operation, write merging, write data forwarding, writing operation, complete parity block write data, weighted histogram operations, load and increment operations, and compare and swap operations. The example main cache store queue **212** is further described below in conjunction with FIG. 4A. The example main cache store queue **212** operates in conjunction with the example main storage **214**. In the example of FIGS. 3A-3D, the main storage **214** is data RAM (DRAM).

(58) The example victim cache store queue **216** of FIGS. 3A-3D includes blocks that correspond to operations of the victim cache store queue **216**. For example, the victim cache store queue **216** includes blocks to implement a read-modify-write operation, write merging, write data forwarding, writing operation, complete parity block write data, load and increment operations, and compare and swap operations. The example victim cache store queue **216** is further described below in conjunction with FIG. 5. The example victim cache store queue **216** operates in conjunction with the example victim storage **218**. In the example of FIGS. 3A-3D, the main storage **214** is a register file.

(59) The example MESI RAM **300** of FIGS. 3A-3D is connected to the command line of the example CPU interface **202** and the example DRAM **214**. The example MESI RAM **300** tracks the state of the data in the example main storage **214** based on the command from the CPU **102**. For example, the MESI RAM **300** tracks the states of a cache line by marking the cache line as modified, exclusive, shared, or invalid. Modified, or dirty, corresponds to when the cache line contains data that is not stored in any other similar-level cache and the data has been modified from the value in main memory (e.g., so that a read of the data in extended memory **106** is flagged or permitted because it is not up-to-date or valid). Exclusive is when the cache line contains data that is not stored in any other similar-level cache and the data is clean (e.g., matches the data in the extended memory **106**). Shared indicates that the cache line contains data that may be stored in

other caches and is clean (e.g., the line may be discarded because it is present in another cache). Invalid indicates that the cache line is invalid or unused. The MESI RAM **300** may be called upon when updates to the main storage **214** and/or the extended memory **106**. The example MESI RAM **300** for victim cache is implemented in conjunction with the example tag RAM **210**.

(60) The example MESI RAM **300** of FIGS. 3A-3D adds coherency to the system by tracking, labelling, marking, etc. The status of a memory address may prevent data writes at particular times so that data mismatches do not occur at different levels of cache. The state tracked by the MESI RAM **300** may be transmitted with cache misses so that the higher level cache understands what the data at the memory address will be used for. For example, if the L1 cache sends out a cache miss for a read operation, the cache miss includes the shared state to let the high level cache know that the data will be read but not manipulated. If the L1 cache **110** sends out a cache miss for a write operation, the cache miss includes an exclusive state to let the higher level cache know that the data will be modified. When the example main storage **214** and/or the victim storage **218** evicts a victim, the status from the example MESI RAM **300** may be used to figure out how to evict. For example, if the data is shared, the main storage **214** and/or victim storage **218** can simply discard (e.g., because the data at memory address already located in a high cache). If the data is exclusive or modified, the cache controller **220** instructs the interface **228** to transmit the victim to the higher level cache (e.g., because the data for the memory address is not located in higher level cache or is located in higher level cache but is outdated).

(61) The example address processing components **302a-c** of FIGS. 3A-3D are connected to the CPU interface **202**, the example main storage **214**, the example main cache store queue **212** (e.g., via the MUX **318**), the example victim storage **218** (e.g., via the example MUX **320**) and each other. The example address processing components **302a-c** include an example first address processing component **302a**, a second address processing component **302b**, and a third address processing component **302c**. The first address processing component **302a** performs address translation, the second address processing component **302b** performs data rotation, and the third address processing component **302c** facilitates bank organization. Alternatively, one or more of the first address processing component **302a**, the second address processing component **302b**, and the third address processing component **302c** can be included in the same hardware, logic circuit, integrated chip, etc. The example address processing components **302a-c** organizes data to be written into the example main storage **214** and the victim storage **218** according to a particular storage protocol to ensure that the data is stored correctly. For example, in a multi-bank example (e.g., where the main cache store queue **212**, the main storage **214**, the victim cache store queue **216**, and the victim storage **218** are broken up into multiple banks), the address processing components **302a-c** may use a memory address from a CPU operation to determine which banks of the main cache store queue **212**, the main storage **214**, the victim cache store queue **216**, and the victim storage **218** is broken up into multiple banks would be needed for the given CPU operation.

(62) The example bank processing logic **303** is coupled to the CPU interface **202**, the example main storage **214**, the example main cache store queue **212** (e.g., via the MUX **318**), and the example victim storage **218** (e.g., via the example MUX **320**). In operation, the bank processing logic **303** is configured to analyze read, modify, and/or write instructions from the CPU interface **202**. In this manner, the bank processing logic **303** is configured to determine the nature of the read, modify, and/or write instructions to facilitate efficient partial bank read, modify, and/or write instructions. In examples disclosed herein, the bank processing logic **303** detects whether incoming write instructions indicate a write of an entire bank, or a write of a partial bank. In this manner, the bank processing logic **303** can indicate whether to operate a read-modify-write operation, while negating to transmit the read instruction. Example description of bank processing logic **303** operation is described below.

(63) The example hit/miss comparison logic **304** of FIGS. 3A-3D is connected to the input CPU interface **202**, the tag RAM **208**, the main storage **214**, the main cache store queue **212**, the cache

controller **220**, and/or the example MUX circuit **314** (e.g., via a data forward latch). The hit/miss comparison logic **304** obtains the address from the tag RAM **208** and an address of the instruction from the CPU **102** and compares the two (e.g., using exclusive nor (XNOR) logic) to determine whether the address from the instruction hit or missed (e.g., the data corresponding to the address is stored in the example DRAM **214** or not). The example hit-miss comparison logic **304** includes TAG compare logic to output the result of the comparison to the example main cache store queue **212**, the example cache controller **220**, and/or to the example MUX circuit **314**.

(64) The example hit/miss comparison logic **306** of FIGS. 3A-3D is connected to the input CPU interface **202**, the tag RAM **210**, the victim cache store queue **216**, and/or the example replacement policy component **308**. The hit/miss comparison logic **306** obtains the entry number of the victim cache (e.g., location) from the tag RAM **210** and an address from the instruction from the CPU interface **202** and compares the two to determine if the access (e.g., the instruction from the CPU interface **202**) is a hit or miss (e.g., the data corresponding to the address is stored in the example victim storage **218** or not). The example hit-miss comparison logic **306** outputs the result to the replacement policy component **308**, the address encoder **326**, the multiplexer **330**, and/or the victim cache store queue **216**.

(65) The address encoder **326** of FIGS. 3A-3D is connected to the tag RAM **210** (via the hit/miss comparison logic **306**) and a multiplexer **330**. The address encoder **326** encodes an address of the tag RAM **210** to a form that is interpretable by the victim storage **218**. For example, the tag RAM **210** may store 16-bit memory addresses while the victim storage **218** stores 4-bit memory addresses corresponding to the 16-bit memory addresses. Thus, the address encoder **326** may transform the 16-bit memory address into a 4-bit memory address to locate and/or enter the corresponding memory address in the victim storage **218**. Additionally and/or alternatively, the address encoder **326** encodes the memory address into any bit value. The address encoder **326** is coupled to an example multiplexer **330**, such that the address encoder **326** provides the encoded address to the multiplexer **330**.

(66) The example replacement policy component **308** of FIGS. 3A-3D is connected to the hit-miss comparison logic **306** and a replacement address encoder **328**. In some examples, the replacement policy component **308** is connected to the tag RAM **210**. The example replacement policy component **308** controls the replacement policy of data stored in the example victim storage **218** (e.g., which data to move out of the victim cache in the event that a new victim is to be stored in the victim storage **218**). In some examples, there may be a replacement policy component for the example main storage **214**. However, a replacement policy is not required for a direct mapped cache (e.g., if the example main storage **214** is directly mapped because a particular address will only go in one location). The example flush engine (e.g., the flush engine component) **309** is coupled to the replacement policy **308**. In some examples, the flush engine **309** is used and/or otherwise invoked to flush out write misses stored inside the victim storage **218** at a pre-defined periodicity.

(67) The replacement address encoder **328** of FIGS. 3A-3D is coupled to the multiplexer **330**, such that the replacement address encoder **328** provides the encoded replacement address to the multiplexer **330**. The replacement address encoder **328** encodes an address selected by the replacement policy component **308**. For example, the replacement policy component **308** outputs a 16-bit replacement memory address to the replacement address encoder **328**, and the replacement address encoder **328** transforms the 16-bit replacement memory address to a 4-bit replacement memory address. Additionally and/or alternatively, the replacement address encoder **328** encodes the replacement memory address into any bit value.

(68) The multiplexer **330** of FIGS. 3A-3D is connected to the hit/miss comparison logic **306**, the address encoder **326**, the replacement address encoder **328**, and an address read **332**. The multiplexer **330** selects an encoded address based on the result of the hit/miss comparison logic **306** (e.g., based on whether the access was a hit or miss). The multiplexer **330** provides the selected

memory address to the address read **332** to be read into the victim storage **218**. For example, the multiplexer **330** outputs the location of the victim storage **218** which the CPU instruction may use to either read the data from or store the data in case of a write instruction.

(69) The example error correcting code (ECC) logic **310** of FIGS. 3A-3D is connected to the example main storage **214** via the latch **322**, and the MUXing circuit **314**. If the example L1 cache **110** supports ECC memory, the example main storage **214** stores data in blocks along with a set of ECC syndrome bits that correspond to the blocks. When a read operation is received, the example main storage **214** may provide the stored data block and the corresponding ECC syndrome bits to the ECC logic **310**. The example ECC logic **310** may regenerate the ECC syndrome bits based on the data block as read from the main storage **214** and compare the regenerated ECC syndrome bits with the previously stored bits. The example ECC logic **310** may determine that a data block has been incorrectly read when there is a discrepancy and may correct the error in the data block.

(70) The example error correcting code (ECC) logic **312** of FIGS. 3A-3D is connected to the example victim storage **218** via the latch **324**, and the MUX circuit **316**. The example ECC logic **312** performs ECC correction logic in a similar manner as the example ECC logic **310**. However, the example ECC logic **312** performs the ECC syndrome bit comparison with respect to the example victim storage **218**.

(71) The example main data MUX circuit **314** of FIGS. 3A-3D is connected to the main cache store queue **212**, the ECC logic **310**, the output of the hit/miss comparison logic **304** (e.g., via a latch), the example output CPU interface **202** (e.g., via the example latch **322**), the example victim data MUX **316**, and the example cache controller **220**. The example main data MUX circuit **314** has the reverse operation of the example address processing components **302a-c** to convert the data in the data address to be transmitted to the example CPU **102**. The MUXs within the example MUX circuit **314** are controlled by the cache controller **220** via the respective select inputs.

(72) The example victim data MUX circuit **316** of FIGS. 3A-3D is connected to the victim cache store queue **216**, the ECC logic **312**, the example main data MUX circuit **314**, the L2 cache **112** (e.g., via a latch), and the example cache controller **220**. The example main data MUX circuit **314** has the reverse operation of the example address processing components **302a-c** to convert the data in the data address to be transmitted to the example CPU **102**. The MUXs within the example MUX circuit **316** are controlled by the cache controller **220** via the respective select inputs.

(73) The example MUX **318** of FIGS. 3A-3D is connected to the L2 cache **112**, the example CPU interface **202** via a latch, the address processing components **302a-c**, the main cache store queue **212**, the main storage **214**, and the cache controller **220**. The example cache controller **220** controls the MUX **318** to control all the ways that data could be written into or read from the example main storage **214** (e.g., via direct memory access (DMA), a snoop transaction (e.g., when the L2 cache **112** wants data from the L1 cache **110**), the main cache store queue **212**, the read hit buffer, etc.). The example MUX **320** is connected to the example main storage **214** (e.g., via the eviction line for storing victims), the address processing components **302a-c** (e.g., via the write-stream and/or CPU load), the example L2 cache **112** (e.g., via the L2 W and/or snoop lines), the victim storage **218**, and/or the cache controller **220**. The example cache controller **220** controls the MUX **320** to control all the ways that data could be written into or read from the example victim storage **218** (e.g., via direct memory access (DMA), a snoop transaction, the victim cache store queue **216**, an eviction from the example main storage **214**, via L2 cache **112**, etc.). The example MUXs **318**, **320** are controlled by the example cache controller **220** via the respective select inputs.

(74) The elements of FIGS. 3A and 3B are arranged in a manner representing the flow of transaction processing. The temporal elements are represented by stages E1-E5. Each stage represents a discrete period of time, and in some examples each stage represents a clock cycle or an integer number of clock cycles.

(75) Write Miss Caching in L1 Data Cache

(76) In operation, the example CPU **102** transmits instructions (e.g., returns data based on a read

operation, write data to a particular memory location). The first time the example CPU **102** requests data for a particular memory address, if the data for the memory address is not stored in the main storage **214**, the tag ram **208** will output a read miss and the L1 cache will send the read miss information out (e.g., to L2 cache **112**, L3 cache **114**, and/or the extended memory **106**) to be serviced (e.g., return the data corresponding to the address of the request from the CPU **102**). The L1 cache **110** may then perform a read allocate. A read allocate is when the L1 cache **110** stores the data in the main storage **214**, updates the tag RAM **208**, etc., to identify that the data for the address is now stored in the main data storage. The L1 cache **110** may return the data to the CPU **102** and/or wait for the CPU **102** to send out a subsequent read request for the same address. If the CPU **102** sends out a subsequent read request for the same address, the tag RAM **208** will identify that the data for the address is now present in the main storage **214**, thereby resulting in a read hit. If the CPU **102** does a write to the same address, the tag RAM **208** will identify a write hit because the address is stored in the main storage **214**. For a write hit, the CPU **102** will provide data to write, and the L1 cache **110** will write the data into the main storage **214** corresponding to the address.

(77) If the initial action from the CPU **102** (or a subsequent action from the CPU **102**) resulted in a write miss (e.g., the memory address that the CPU **102** is attempting to write to is not stored in the example main storage), the L1 cache **110** can perform a write miss. During a write miss, the L1 cache **110** sends the write miss out to the higher level cache (e.g., L2 cache **112**, L3 cache **114**, etc.) and/or extended memory **106** to retrieve the data from the memory address, stores the data in the main storage **214**, and then writes the data from the CPU **102** in the main storage **214** at a location corresponding to the memory address.

(78) However, when the CPU **102** performs a write operation, it typically does not need the information that it has written right away. Accordingly, some example cache systems included a write buffer to store write instructions from the CPU to prevent stalling the CPU (e.g., by attempting to load memory address from higher cache before writing the data). In this manner, the CPU can continue to send instructions while the L1 cache interfaces with higher level cache to obtain the data in the memory address corresponding to the write instructions in the buffer, and stores the data in the main storage and writes the data (e.g., stored in the write buffer) in parallel with subsequent instructions from the CPU. However, the CPU may only write a few number of bytes per write instruction and the interface between the L1 cache and higher level caches and/or the extended memory is capable of sending a larger number of bytes (e.g., 64 byte bandwidth). Accordingly, the transmission of a few number of bytes per cycle on a large byte interface is inefficient.

(79) In the system of FIGS. **3A-3D**, instead of implementing a write buffer, a section of the victim storage **218** is dedicated for write-miss information. Accordingly, the example victim storage **218** is a victim cache and a write miss buffer. The section of the victim storage is called the write miss cache. In some examples, the write miss cache may be additionally or alternatively implemented in the main storage **214**. In some examples, the write miss cache is a 128 bytes of a cache line. The write miss cache stores all the write miss data until the write miss cache is full and/or there is more than a first threshold number of bytes that can be sent to higher level cache and/or extended memory. Once the write miss cache is full or the first threshold amount of write miss data has been added to the write miss cache, the victim storage **218** combines a second threshold amount of the write miss data in the write miss cache into one signal that is sent to the higher level cache (e.g., via the example L2 interface **228**) to be written in the address stored in the higher level cache (e.g., the L2 cache **112**) and/or the extended memory **106**. In this manner most or all of the bandwidth of the interface can be utilized in a particular cycle. The second threshold may be the same as or different than the first threshold. Additionally, the write data is stored locally in the main storage **214** or the victim storage **218**. Thus, if an additional read or write operation is sent from the CPU **102** prior to the write operation being passed to the higher cache (e.g., the L2 cache **112**, the example L3 cache **114**, etc.) and/or the extended memory **106**, the cache controller **220** can read and/or write the data

to the corresponding address within the write miss cache before it gets transmitted to higher level cache and/or the extended memory **106**.

(80) In some examples, the structure of the write miss cache in the victim storage **218** includes a byte enable register file that represents the value bytes (e.g., the bytes to be written) of the write miss information. For example, if a write miss corresponding to writing data for a first byte and a third byte of a memory address is stored in the write miss cache, the victim storage **218** stores the write miss data for the first and third byte in conjunction with the memory address and populates the corresponding entry of byte enable register file with a first value (e.g., '1') for the elements of the entry that correspond to the first and third byte and a second value (e.g., '0') for the remaining elements of the entry. In this manner, when the write miss data is sent to higher level cache, the byte enable bits of the entry are included in the transmission so that the higher level cache knows which data is valid (e.g., which bytes are to be written to) and which data is invalid (e.g., which bytes should not be written to).

(81) In certain cases, if a read or a write request is a hit in the main storage **214**, the results of the victim storage **218** may be disregarded. Where the main storage **214** is configured as a one way associative cache, cache request conflicts as with multiple datapaths are straight forward. A datapath is one or more logic circuits that, when executed and/or otherwise invoked, effectuate the transmission or delivery of data (e.g., logic signal(s), bit vector(s), etc.) from a source to a destination. Cache request conflicts from the multiple datapaths may be an issue for the victim storage **218** where the victim storage **218** is fully associative and can store write-miss entries as well as cache entries. Cache request conflicts for two datapaths may be handled as described in the following examples. Of note, in the following examples, cache operations are described in the context of a particular datapath. However, it may be understood that the examples are directed to combinations of particular cache operations, and which datapath a particular cache operation is on does not matter.

(82) In a first example, there is a cache read miss for a first datapath and a cache write hit for a second datapath, with two different memory addresses. In certain cases where the victim storage **218** is fully associative, cache requests can go to any location within the victim storage **218**. Referring to FIGS. 3A-D, address generation for a location within the victim storage **218** occurs in the E2 pipestage along with address lookup in the victim cache tag RAM **210**, thus address generation for a location within the victim cache is performed before it is known whether an address of the cache request is a hit or miss. After the controller **220** determines that there is a read miss on the main storage **214**, a request to obtain the memory address of the read request is issued to a higher level cache or memory. Where the main storage **214** is one way associative, the memory address of the read request of the first datapath maps to a single location in the main storage **214**. If there is already cached data in the single location, the already cached data is evicted from the main storage **214** to the victim storage **218** to the pre-generated location within the victim storage **218**. If this pre-generated location is the same location the cache write of the second datapath is a hit on, a conflict occurs. This conflict may be detected by the cache controller **220**. As indicated above, address generation for a location within the victim storage **218** occurs before it is known whether the address of cache request is a hit or a miss, thus there is an address generated for a second location within the victim storage **218** for the cache write of the second datapath prior to the determination that the cache write is a hit. Based on the detection of the conflict, this second location within the victim cache may be used to store the data evicted from the main storage **214** by the read miss.

(83) In another example, there may be a cache read miss for the first datapath and a cache write hit for the second datapath, as with the first example. In this example, due to the read miss, a request to obtain the memory address of the read request is issued to a higher level cache or memory and the already cached data is evicted from the main storage **214** to the victim storage **218** to a pre-generated location, here location A, within the victim storage **218**. In this example, the cache write

of the second datapath hits on location A within the victim storage **218** as well, resulting in a set conflict. One possible solution to such a conflict is to load the requested read miss from the higher level cache or memory directly to the victim cache in a separate location. Another solution to the conflict is for the cache controller **220** to stall the read miss such that the already cached data is not evicted to the victim storage **218** until after the cache write of the second datapath completes to location A and location A is evicted to a higher level cache or memory. In certain cases, the stall may occur while waiting for the higher level cache or memory to return the read miss data.

(84) As another example, a first cache write may be received on the first datapath for the victim storage **218** and a second cache write is also received on the second datapath for the victim storage **218**. If the first cache write and the second cache write are addressed to different memory addresses, then both cache writes may proceed in parallel to the victim storage **218**. If the first cache write and the second cache write are both addressed to the same memory address and both are cache misses, then the victim cache controller **224** allocates a single location in the allocated victim cache write-miss memory and merges the first cache write and the second cache write in the victim cache store queue **216**. The merged data may then be written to the victim cache write-miss memory along with merged byte enable, priority, and color tag information. If the first cache write and the second cache write are both addressed to the same memory address and both are cache hits, both the first cache write and the second cache write are merged into the cache at the hit location.

(85) As another example a cache read may be received on the first datapath for the victim storage **218** and a cache write may be received on the second datapath for the victim storage **218**. In the cases where the cache read is a hit and the cache write is a hit to different locations in the victim storage **218**, the cache read and cache write proceed in parallel without conflicts. In the case the cache read is a hit and the cache write is a miss to different locations in the victim storage **218**, the cache read and cache write also proceed in parallel without conflicts. In the case the cache read is a miss and the cache write is a hit to different locations in the victim storage **218**, the cache read may use an address generated for a location within the victim storage **218** for the cache write as discussed above. In the case the cache read is a miss and the cache write is a miss to different locations in the victim storage **218**, both the cache read and the cache write use addresses generated for locations within the victim storage **218**.

(86) In another example, a cache read may be received on the first datapath for the victim storage **218** and a cache write may be received on the second datapath for the victim storage **218**. In certain cases, the cache read may be a hit for a first address of a set of addresses stored in the victim storage **218**. The cache write may also be a hit for a second address of the same set of addresses stored in the victim storage **218**. In such a case, the cache read and the cache write may proceed in parallel without conflicts. In another case, the cache read may be a hit for a first address of a set of addresses stored in the victim storage **218**. The cache write may be a miss for a second address of the same set of addresses stored in the victim storage **218**. In such a case, the cache read and the cache write may proceed in parallel without conflicts. In another case, the cache read may be a miss for a first address of a set of addresses stored in the victim storage **218**. The cache write may be a hit for a second address of the same set of addresses stored in the victim storage **218**. In such a case, as similar to that discussed above, the cache read may be stalled until after the cache write of the second datapath completes to the location in the victim storage **218** and is evicted to a higher level cache or memory. The cache read then proceeds to read the set of addresses from the higher level cache or memory into the victim storage **218**. In another case, the cache read may be a miss for a first address of a set of addresses stored in the victim storage **218**. The cache write may also be a miss for a second address of the same set of addresses stored in the victim storage **218**. In such a case, the cache read and the cache write may proceed in parallel without conflicts.

(87) In another example, a cache read may be received on the first datapath for the victim storage **218** and a cache write may be received on the second datapath for the victim storage **218**. In certain cases, the cache read may be a hit for an address stored in the victim storage **218**. The cache write

may also be a hit for the same address stored in the victim storage **218**. In such a case, the cache read may proceed first and the cache write may be stalled until after the cache read completes. Alternatively, the order of the cache write and cache read may be based on the datapath on which the cache write and cache read are received, with the cache operation arriving on a lower (or higher) numbered datapath being completed before the other cache operation. In certain cases, the cache read may be a miss for an address stored in the victim storage **218**. The cache write may also be a miss for the same address stored in the victim storage **218**. In such a case, the cache write operation may be forwarded to a higher level cache or memory and then the cache read may obtain the data from the higher level cache or memory after the cache write operation completes for storage into the victim storage **218**.

(88) In another example, a first cache read may be received on the first datapath for the victim storage **218** and a second cache read may be received on the second datapath for the victim storage **218**. If the first cache read and the second cache read are for different memory addresses, then there are no conflicts for either hits nor misses. In certain cases, the first cache read may be a miss for a first address of a set of addresses. The second cache read may also be a miss for a second address of the same set of addresses. If the first cache read and the second cache read have different priority levels, a higher level cache or memory is accessed based on the higher of the different priority levels. Otherwise, the higher level cache or memory is accessed and the set of memory addresses obtained for storage in the victim storage **218**. The case where the first cache read and the second cache read are for the same address is handled identically.

(89) Example methods, apparatus, systems, and articles of manufacture to facilitate Write Miss Caching in L1 data cache are disclosed herein. Further examples and combinations thereof include the following: Example 1 includes an apparatus comprising a first cache storage, a second cache storage, wherein the second cache storage includes a first portion operable to store a first set of data evicted from the first cache storage and a second portion, a cache controller coupled to the first cache storage and the second cache storage and operable to receive a write operation, determine that the write operation produces a miss in the first cache storage, and in response to the miss in the first cache storage, provide write miss information associated with the write operation to the second cache storage for storing in the second portion.

(90) Example 2 includes the apparatus of example 1, wherein the cache controller is operable to compare the second portion of the second cache storage to a threshold, and based on the second portion exceeding the threshold, cause the write miss information to be transmitted to a second cache.

(91) Example 3 includes the apparatus of example 2, wherein the threshold corresponds to a bandwidth of an interface.

(92) Example 4 includes the apparatus of example 2, wherein the threshold corresponds to a size of the second portion.

(93) Example 5 includes the apparatus of example 1, wherein the write miss information is first write miss information, the cache controller operable to the second cache storage is to not provide the first write miss information from the first cache storage to the second cache storage if the second cache storage includes second write miss information for a memory address corresponding to the first write miss information.

(94) Example 6 includes the apparatus of example 1, wherein the first cache storage and the second cache storage are connected in parallel to a central processing unit.

(95) Example 7 includes the apparatus of example 1, wherein the write miss information is first write miss information, the cache controller operable to, when a first memory address of the first write miss information from the first cache storage matches a second memory address of second write miss information stored in the second portion, to merge the first write miss information with the second write miss information.

(96) Example 8 includes the apparatus of example 7, wherein the cache controller is to merge the

first write miss information with the second write miss information by at least one of (a) maintaining first write information of the first write miss information or (b) discarding second write information of the second write miss information when the second write information corresponds to the same one or more bytes as the first write miss information.

(97) Example 9 includes the apparatus of example 1, wherein the second portion includes a byte enable register, the cache controller to store values in the byte enable register based on the write miss information.

(98) Example 10 includes the apparatus of example 9, wherein the values correspond to elements of the write miss information that will be written.

(99) Example 11 includes a system comprising a central processing unit to output a write command corresponding to a memory address, a first cache storage to output write miss information from a first storage to a second storage when the first cache storage does not store data for the memory address, and a second cache storage, wherein the second cache storage includes a first portion operable to store a first set of data evicted from the first cache storage and a second portion, the second cache storage to store the write miss information in a dedicated section of the second storage, the dedicated section being dedicated to the write miss information.

(100) Example 12 includes the system of example 11, wherein the second cache storage is to, when the dedicated section has more than a threshold amount of write miss information, output the write miss information to a second cache.

(101) Example 13 includes the system of example 12, wherein the threshold corresponds to a bandwidth of an interface to the second cache.

(102) Example 14 includes the system of example 12, wherein the threshold corresponds to a size of the second portion.

(103) Example 15 includes the system of example 11, wherein the second cache storage is to not store the write miss information from the first storage in the second portion if the second cache storage includes second write instructions corresponding to a same memory address as the write miss information from the central processing unit.

(104) Example 16 includes the system of example 11, wherein the first cache storage and the second cache storage are connected in parallel to the central processing unit.

(105) Example 17 includes the system of example 11, wherein the write miss information is first write miss information, further including a controller to, when a first memory address of the first write miss information from the first cache storage matches a second memory address of second write miss information stored in the second portion, merge the first write miss information with the second write miss information.

(106) Example 18 includes the system of example 17, wherein the cache controller is to merge the first write miss information with the second write miss information by at least one of (a) maintaining first write information of the first write miss information or (b) discarding second write information of the second write miss information when the second write information corresponds to the same one or more bytes as the first write miss information.

(107) Example 19 includes a method comprising receiving a write operation, determining that the write operation produces a miss in a first cache storage, and in response to the miss in the first cache storage, providing write miss information associated with the write operation to a second cache storage element with a first portion and a second portion for storing in the second portion, the first portion storing a first set of data evicted from the first cache storage.

(108) Example 20 includes the method of example 19, further including comparing the second portion of the second cache storage to a threshold, and based on the second portion exceeding the threshold, outputting the write miss information to a second cache.

(109) FIG. 4A is an example circuit implementation of the main cache store queue **212** of FIGS. 2 and/or 3. In FIG. 4A, the main cache store queue **212** includes an example latches **402a**, **402b**, **402c**, **402d**, **402e**, example merge circuits **403a-c**, an example arithmetic component **404**, an

example atomic compare component **406**, an example read-modify-write merge component **408**, an example select multiplexer **410**, and example ECC generator **412**, an example arbitration manager **414**, an example pending store address data store **416**, an example priority multiplexer **418**, an example read port **424**, and an example write port **426**. The example merge circuits **403a-d** include an example comparator(s) **420**, and example switches **422**. The example of FIG. 4A illustrates a single pipeline of the main cache store queue **212**. However, the main storage element **214** may be arranged to support more than one independent copy of the pipeline with respect to different banks as indicated by the dashed box **400**. Accordingly, the pipeline of FIG. 4A may be reproduced multiple times for different banks, as further described below.

(110) In FIG. 4A, the example latches **402a**, **402b**, **402c**, **402d**, **402e**, are electronic devices configured to store information (e.g., bytes, bits, etc.) obtained by the main cache store queue **212**. The example latches **402a-c** pass the write data and information corresponding to whether the write data needs to be combined with the read and corrected data out of the ECC logic **310** in the arithmetic unit **404**, the atomic unit **406**, and/or the RMW merge component **408**. In the example of FIG. 4A, the latch **402a** is communicatively coupled to the cache controller **220** to obtain read, write, and/or modify instructions. Such read, modify, and/or write instructions may originate from the CPU **102**, and transmitted to the latch **402a** via the cache controller **220**. Latch **402a** is coupled to latch **402b**, the tag RAM **208**, the arbitration manager **414**, and the pending store address data store **416** to transmit such read, modify, and/or write instructions to the latch **402b**, the tag RAM **208**, the arbitration manager **414**, and the pending store address data store **416** in response to a subsequent clock cycle of the cache controller **220**.

(111) In the example of FIG. 4A, the latch **402b** is coupled to latch **402a**, the pending store address data store **416**, latch **402c**, the priority multiplexer **418**, and the comparator(s) **420** to transmit the obtained read, modify, and/or write instruction in response to a subsequent clock cycle of the cache controller **220**. In this manner, the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.), propagates through the main cache store queue **212**.

(112) The example latch **402c** is coupled to the latch **402b**, the priority multiplexer **418**, the arithmetic component **404**, the atomic compare component **406**, and the read-modify-write merge component **408**. This coupling enables the latch **402c** to transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to the arithmetic component **404**, the atomic compare component **406**, and/or the read-modify-write merge component **408** in response to a subsequent clock cycle of the cache controller **220**. In some examples disclosed herein, latch **402c** may transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to one or more of the arithmetic component **404**, the atomic compare component **406**, and/or the read-modify-write merge component **408**. For example, in the event the cache controller **220** transmits an instruction to write a bit into a currently stored word, the latch **402c** may transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to the read-modify-write merge component **408**. In other examples disclosed herein, though the cache controller **220** may transmit an instruction to write a bit into a currently stored word, the latch **402c** may transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to the arithmetic component **404**, the atomic compare component **406**, and/or read-modify-write merge component **408**.

(113) The example latch **402d** is coupled to the first multiplexer **410**, the ECC generator **412**, and the latch **402e**. This coupling enables the latch **402d** to transmit the value obtained from first multiplexer **410** to ECC generator **412** and/or the merging circuit **403a** in response to a subsequent clock cycle of the cache controller clock **220**.

(114) The example latch **402e** is coupled to priority multiplexer **418** and to the MUX circuit **314** of FIGS. 3 and/or 4. This coupling enables the latch **402e** to transmit the value obtained from the from the priority multiplexer **418** to MUX circuit **314** in response to a subsequent clock cycle of the

cache controller **220**. For example, on a read operation by the CPU **102** to an address that hits the main storage element **214**, a previous write to the same address could still be inflight (e.g., not completely written to the main storage element **214**). In such an example, the latch **402e** collects the unwritten data byte and provides it to the MUX **314** which eventually gets sent back to the CPU **102**.

(115) The example merging circuit **403a** is coupled to the latch **402d**, the merging circuit **403b**, the arithmetic component **404**, the atomic compare component **406**, and the read-modify-write merge component **408**. The example merging circuit **403b** is coupled to the merging circuit **403a**, the priority multiplexer **418**, and the merging circuit **403c**. The example merging circuit **403c** is coupled to the merging circuit **403b** and the latch **402b**. The example merging circuits **403a-c** facilitate the comparison of read operations in different sections of the main cache store queue **212** to potentially reroute write operations to be merged with write operations corresponding to the same memory address location, as further described below. Although the example of FIG. **4A** includes three merging circuits **403a-c**, there may be additional merging circuits to merge write operations from other sections of the main cache store queue **212** (e.g., a merging circuit coupling the output of the latch **402d** to the output of latch **402b** and/or latch **402a**, etc.). In some examples, the merging circuits **403a-c** are combined into a single circuit that compares the write operations from the different latches **402b-d** and reroutes based on matching memory addresses in any two or more of the different latches **402b-d**.

(116) In the example illustrated in FIG. **4A**, the arithmetic component **404** is coupled to the latch **402c**, the first multiplexer **410**, and to the ECC logic **310** to perform arithmetic operations on (e.g., increment, decrement, etc.) data from the main storage **214**. Additionally, the arithmetic component **404** performs histogram operations on the data stored in the main storage **214**. The example arithmetic component **404** of the illustrated example of FIG. **4A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. Operation of the example arithmetic component **404** is further described below.

(117) In the example illustrated in FIG. **4A**, the atomic compare component **406** is coupled to the latch **402c**, the first multiplexer **410**, and to the ECC logic **310** to compare data at a memory address to a key and, in the event the data at the memory address matches the key, replace the data. The example atomic compare component **406** the illustrated example of FIG. **4A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. Operation of the example atomic compare component **406** is further described below.

(118) In the example illustrated in FIG. **4A**, the read-modify-write merge component **408** is coupled to the latch **402c**, the first multiplexer **410**, and to the ECC logic **310** to facilitate the read, modify, and/or write instruction(s) sent by the cache controller **220**. For example, the read-modify-write merge component **408** is coupled to the ECC logic **310** to obtain the currently stored word that is to be affected by the read, modify, and/or write instruction(s). In an example write operation, the read-modify-write merge component **408** is configured to update the currently stored word obtained from the ECC logic **310** with the new bit(s), byte(s), etc., obtained from the latch **402c**. Additional description of the read-modify-write merge component **408** is described below. The example read-modify-write merge component **408** of the illustrated example of FIG. **4A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog

or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(119) The example first multiplexer **410** is coupled to the arithmetic component **404**, the atomic compare component **406**, and the read-modify-write merge component **408** to transmit, based on an indication from the cache controller **220**, the output of either the arithmetic component **404**, the atomic compare component **406**, or the read-modify-write merge component **408** to the latch **402d**. For example, in the event the cache controller **220** indicates to perform a write function (e.g., the cache control transmits a write request to the latch **402b**), an indication is sent by the cache controller **220** to the first multiplexer **410** to select the input connected to the read-modify-write merge component **408** to be transmitted to the latch **402d**. The example first multiplexer **410** of the illustrated example of FIG. **4A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(120) In the example illustrated in FIG. **4A**, the ECC generator **412** is coupled to the latch **402d** and to the merging circuit **403a** to facilitate error detection and correction in the value (e.g., byte(s), bit(s), etc.) stored in the latch **402d**. For example, the ECC generator **412** is configured to regenerate the ECC (E.g., generate error detection code) value which will be stored with the data (e.g., merged word output from the read-modify-write merge component **1108**). The ECC value is used by the error detection and correction circuit to determine whether the error occurred during a read and/or write operation, as further described above. The example ECC generator **412** of the illustrated example of FIG. **4A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(121) In FIG. **4A**, the example arbitration manager **414** is coupled to the latch **402a**, the latch **402b**, the pending store address datastore **416**, and the main storage **214** to facilitate the read, modify, and/or write instructions obtained from the cache controller **220**. For example, in the event a write instruction is received from the cache controller **220**, the arbitration manager **414** is configured to transmit a read instruction of the corresponding currently stored word to the main storage **214**. In addition, the arbitration manager **414** is coupled to the main storage **214** to arbitrate between conflicting accesses of the main storage **214**. When multiple operations attempt to access the main storage **214** in the same cycle, the arbitration manager **414** may select which operation(s) are permitted to access the main storage **214** according to a priority scheme. Suitable priority schemes are described in more detail in below; however, in one example, the arbitration prioritizes read operations over write operations because write data that is in the main cache store queue **212** is available for use by subsequent operations even before it is written to the main storage **214**. Thus, there is minimal performance impact in allowing the write data to wait in the main cache store queue **212**. However, as the main cache store queue **214** fills with write data that has not yet been written back, the priority of the write operations may increase until they are prioritized over competing read operations.

(122) The example arbitration manager **414** of the illustrated example of FIG. **4A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(123) In the example of FIG. 4A, the pending store address data store **416** is configured to store the address of the read, modify, and/or write instruction obtained from the cache controller **220**. In this manner, the pending store address data store **416** maintains a log of the addresses associated with each value stored in any of the latches **402a**, **402b**, **402c**, **402d**, **402e**, and/or the merging circuits **403a**, **403b**, and/or **403c**. The example pending store address data store **416** of the illustrated example of FIG. 4A may be implemented by any device for storing data such as, for example, flash memory, magnetic media, optical media, etc. Furthermore, the data stored in the pending store address data store **416** may be in any data format such as, for example, binary data, comma delimited data, tab delimited data, structured query language (SQL) structures, etc.

(124) The example priority multiplexer **418** is coupled to the latch **402b**, the latch **402c**, the latch **402d**, and the merging circuit **403a** to facilitate read operations in the event either of the of the latch **402b**, the latch **402c**, the latch **402d**, or the merging circuit **403a** are storing a value corresponding to a write instruction. For example, the cache controller may initiate the following four write instructions regarding a four byte word having addresses A3, A2, A1, and A0: write address A0 with the byte 0x11, write address A1 with the byte 0x22, write address A3 with the byte 0x23, and write address A0 with the byte 0x44. In such an example, in the event the cache controller **220** initiates a read instruction to read the entire four byte word, the priority multiplexer **418** is configured to obtain the byte value 0x11 stored in the merging circuit **403a**, the byte value 0x22 stored in the latch **402d**, the byte value 0x23 stored in the latch **402c**, and the byte value 0x22 stored in the latch **402b**. Additionally, the pending store address data store **416** transmits an instruction to the priority multiplexer **418** indicating which address value is associated with the byte value stored in the latch **402b**, the latch **402c**, the latch **402d**, and the merging circuit **403a**. In such an example, since the byte value stored in the latch **402b** is associated with the address A0 and the byte value stored in the merging circuit **403a** is also associated with the address A0, the priority multiplexer **418** is configured to transmit a packet to the latch **402e** indicating that address A0 is 0x44 (e.g., the most recent write instruction associated with the address A0), address A1 is 0x22, and address A3 is 0x23.

(125) In this manner, the MUX circuit **314** is configured to update the value of the currently stored word with the byte values obtained from the priority multiplexer **418**. Such an operation ensures that a read instruction transmitted by the main cache store queue **212** probably indicates the correct word, even though the write instructions may not have fully propagated through the main cache store queue **212**.

(126) An example read path (e.g., the read input to the tag RAM **208**) may run in parallel with the main cache store queue **212**. Because a read operation (e.g., a read instruction) may refer to data in a write operation (e.g., a write instruction) that may not have completed yet, the main cache store queue **212** may include write forwarding functionality that allows the read path to obtain data from the main cache store queue **212** that has not yet been written back to the main storage **214**. In an example, the main cache store queue **212** includes pending store address data store **416** that records the addresses of the operations at each stage of the main cache store queue **212**, a priority multiplexer **418** to select data from one of the stages (e.g., latches) of the main cache store queue **212** for forwarding, and a MUX circuit **314** that selects between the output of the main storage **214** (by way of the error detection and correction circuit **310**) and the forwarded main cache store queue **212** data from the data priority multiplexer **418**.

(127) Alternatively, because a read operation (e.g., a read instruction, a read command, etc.) may refer to victimizing (e.g., evicting) data in a write operation (e.g., a write instruction, a write command, etc.) that may not have completed yet, the main cache store queue **212** includes a read-invalidate functionality that forwards in-flight data (e.g., data of the store queue **212** not yet stored in the main storage element **214**) to the victim storage element **214** and/or the L2 cache **212** and invalidates the in-flight data remaining in the store queue **212**.

(128) The example read port **424** is coupled to the read path and the data store **416**. The read port

424 may be implemented by an interface that interfaces with the main cache controller **222** whenever a read-miss occurs. For example, the read port **424** is utilized to receive victim addresses and read-invalidate commands from the main cache controller **222**. The read port **424** is to send the victim addresses to the data store **416** to be compared against the pending addresses stored in the data store **416**.

(129) The example write port **426** is coupled to the write path and the latch **402a**. The write port **426** may be implemented by an interface that interfaces with the main cache controller **222** (e.g., the cache controller **220**) to obtain a write instruction. For example, the write port **426** is utilized to receive addresses and values from the cache controller **220** to write.

(130) The elements of FIG. **4A** are arranged in a manner representing the flow of transaction processing. The temporal elements are represented by stages E1-E5. Each stage represents a discrete period of time, and in some examples each stage represents a clock cycle or an integer number of clock cycles.

(131) An example flow of a read operation through the main cache store queue **212** will now be described. In a first cycle, corresponding to stage E1, the L1 data cache **110** retrieves a record from the tag RAM **208** that is associated with an address of the read operation to determine whether the data is stored in the main storage **214**. In a direct mapped example, the L1 data cache **110** need not wait for the tag RAM **208** comparison before requesting data from the main storage **214**, and thus, the tag RAM **208** comparison between the address of the read operation and the record of cached addresses may extend into a second or third clock cycle.

(132) In the second cycle, corresponding to stage E2, the L1 data cache **110** may request the data and ECC syndrome bits from the main storage **214** if the arbitration manager **414** permits. In this cycle, the L1 data cache **110** may also determine whether newer data is available in the main cache store queue **212** by comparing the read address to the pending store address data store **416**. If so, the priority multiplexer **418** is set to forward the appropriate data from the main cache store queue **212**.

(133) Data and ECC may be provided by the main cache storage **214** in the third cycle, corresponding to stage E3. However, this data may or may not correspond to the memory address specified by the read operation because the L1 data cache **110** may allocate multiple extended memory addresses to the same entry in the cache's main storage **214**. Accordingly, in the third cycle, the L1 data cache **110** determines whether the provided data and ECC from the main storage **214** corresponds to the memory address in the read operation (e.g., a cache hit) based on the comparison of the tag RAM **208** record. In the event of a cache hit, the data and ECC bits are received by the error detection and correction circuit **310**, which corrects any errors in the data in a fourth cycle, corresponding to stage E4.

(134) As explained above, newer data that has not yet been written to the main storage **214** may be present in the main cache store queue **212**, and may be forwarded from the main cache store queue **212** by the priority multiplexer **418**. If so, the MUX circuit **314** selects the forwarded data over the corrected data from the main storage **214**.

(135) Either the corrected data from the main storage **214** or the forwarded data from the main cache store queue **212** is provided to the L1 data cache **110** in a fifth cycle, corresponding to stage E5. In this way, the controller **220** may provide data with full ECC checking and correction in the event of a cache hit in about 5 cycles.

(136) In the event that the data and ECC bits are not present in the main storage **214** (e.g., a cache miss), the main cache store queue **212** may stall until the data can be retrieved from either the extended memory **106** and/or the victim storage **218**, at which point the data may be written to the main storage and the tag RAM **208** may be updated.

(137) Methods and Apparatus to Facilitate Fully Pipelined Read-Modify-Write Support in Level 1 Data Cache Using Store Queue and Data Forwarding

(138) In examples disclosed herein, there are two cache systems operating in tandem (e.g., the main

storage **214** and the victim cache storage **218**). In examples disclosed herein, the main storage **214** is a direct mapped cache element and the victim cache storage **218** is a fully associative cache storage. Both the direct mapped main storage **214** and the fully associative victim cache storage **218** are protected by an error correcting code (ECC). With this, example applications include reading a line from main storage **214** and/or moving a line from the main storage **214**. As such, ECC logic would correct the cache line and regenerate the ECC syndrome, and then write the line to victim cache storage **218**. Such an application may utilize two additional clock cycles of the CPU (e.g., one clock cycle for error correction by the ECC logic and another clock cycle for ECC syndrome regeneration). To reduce this latency, examples disclosed herein include utilizing the same parity block size between the main storage **214** and the victim cache storage **218**. Thus, both the main storage **214** and the victim cache storage **218** calculate and/or otherwise determine parity on a 32-bit boundary. In this manner, the L1 cache **110** can move a cache line directly from main storage **214** to the victim cache storage **218** with less latency.

(139) In examples disclosed herein, the L1 data cache **110** supports a number of operations that read data from the cache and make changes to the data before rewriting it. For example, the L1 data cache **110** may support read-modify-write operations. A read-modify-write operation reads existing data and overwrites at least a portion of the data. In ECC embodiments, a read-modify-write operation may be performed when writing less than a full bank width (e.g., when writing less than a 32-bit, 4-byte word). The read functionality of the read-modify-write is used because the portion of the data in the bank that will not be overwritten still contributes to the ECC syndrome bits.

(140) A read-modify-write operation may be split into a write operation and a read operation, and the main cache store queue **212** may be structured such that the read operation in the read path stays synchronized with the write operation in the main cache store queue **212**. The read operation and the write operation remain synchronized until the read-modify-write merge component **408** overwrites at least a portion of the read data with the write data to produce merged data. The merged data is provided to the ECC generator **412** that generates new ECC syndrome bits for the merged data, and then the merged data and ECC syndrome bits may be provided to the arbitration manager **414** for storing in the main storage **214**.

(141) In an example read-modify-write operation, the cache controller **220** of FIGS. 2 and/or 3 transmits a write request indicating a subset of bytes of a word, or an entire word (e.g., a first set of data), to be written. In such an example, the write request transmitted by the cache controller **220** includes an address value of the byte(s) to be written and the data to write to the indicated bytes. For example, the main storage **214** of FIGS. 2 and/or 3 may include the four-byte word 0x12345678 associated with addresses A3, A2, A1, A0. In such an example, address A3 corresponds to the byte 0x12, address A2 corresponds to the byte 0x34, address A1 corresponds to the byte 0x56, and address A0 corresponds to the byte 0x78 of the stored word. In such an example operation, the cache controller **220** may transmit a write request to replace address A3 with the byte 0x33, replace address A1 with the byte 0x22, and replace address A0 with the byte 0x11 of the currently stored word 12345678. Further in such an example, the first write request to replace address A3 of the stored word with the byte 0x33 would result in the stored word becoming 0x33345678, the second write request to replace address A1 of the stored word with the byte 0x22 would result in the stored word becoming 0x33342278, and the third write request to replace address A0 of the stored word with the byte 0x11 would result in the stored word becoming 0x33342211.

(142) To process such a request, the L1 data cache **110** compares an address of the read operation to a record of the tag RAM **208**, and the L1 data cache requests the data and ECC syndrome bits from the main storage **214** and/or the main cache store queue **212**. Because the read-modify-write operation will modify the data, in examples that track MESI (Modified, Exclusive, Shared, and Invalid) states of entries in the main storage **214**, a cache hit that is not in either the Modified or Exclusive state may be considered a cache miss. When the data is obtained in the proper state and

any errors are corrected, it is provided to the read-modify-write merge component **408**. In this same cycle, the read-modify-write merge component **408** may overwrite at least a portion of the corrected data with the write data to produce merged data. The ECC generator **412** generates new ECC syndrome bits for the merged data.

(143) In examples disclosed herein, the ECC generator **412** operates on word granularity. That is, the ECC generator **412** calculates the ECC syndrome for a block of data. In examples disclosed herein, the block of data may four bytes (e.g., a word). The main cache store queue **212** processes the write instruction by, at a first example cycle (e.g., to complete the first write request of replacing address A3 with the byte 0x33), because the ECC generator **412** operates on word granularity (e.g., a 4-byte or 32-bit word), the cache controller **220** initiates a read request of the currently stored byte in address A3 of the currently stored word. In this example, the byte and address in the first write request (e.g., 0x33 and A3) is stored in the latch **402b**. In addition, the cache controller **220** transmits a read request of the entire currently stored word to the main storage **214**. Thus, at the end of the first cycle, a read request of the entire currently stored word is transmitted to the main storage **214** and the byte 0x33 is stored in the first latch **402b**.

(144) At an example second cycle (e.g., to complete the second write request of replacing address A1 with the byte 0x22) the byte from the first write request is transmitted to the latch **402c**, the entire currently stored word is transmitted from the main storage **214** to the ECC logic **310**, and the second write request (e.g., to replace address A1 with the byte 0x22) is transmitted by the cache controller **220** to be stored in the latch **402b**. During the second cycle, the read-modify-write merge component **408** obtains the byte stored in the latch **402c** and the entire currently stored word transmitted by the ECC logic **310**. In this manner, the read-modify-write merge component **408** identifies the address of the byte in the currently stored word to be updated. Once the read-modify-write merge component **408** identifies and/or otherwise obtains (a) the value (e.g., byte value, bit value, etc.) of the portion of the currently stored word to be updated from the latch **402c** and the (b) currently stored word from the ECC logic **310**, the read-modify-write merge component **408** writes (e.g., replaces, merges, etc.) the portion of the currently stored word with the value of the portion of the currently stored word obtained from the latch **402c**. For example, the read-modify-write merge component **408** writes the value of the portion of the word to an address value corresponding to the portion of the word in the word. Such an example written portion output by the read-modify-write merge component **408** may be referred to herein as the merged word. In some examples disclosed herein, such a merged word is provided by the read-modify-write merge component **1108** for writing to the victim storage **218**.

(145) In response, to initiate the write portion of the example read-modify-write operation, the select multiplexer **410** transmits the merged word from the read-modify-write merge component **408** to be stored in the latch **402d**. At a subsequent clock cycle, the ECC generator **412** obtains the merged word from the latch **402d** and generates the corresponding ECC syndrome bits. At either the same, or a subsequent clock cycle, the ECC generator **412** transmits the merged word through the merging circuits **403a**, **403b**, and **403c** to be handled by the arbitration manager **414** to be stored in the main storage **214**.

(146) However, in such an example described above, because three consecutive write instructions were transmitted by the cache controller **220** corresponding to the same stored word in the main storage **214**, the read-modify-write merge component **408** can merge these three write instructions into one transaction when writing the three values of the three write instructions to the currently stored word. For example, prior to storing the merged word from a first write instruction in the main storage **212**, the main storage queue **212** feedbacks (e.g., transmits) the merged word for use in the subsequent second write instruction.

(147) FIG. 4B is an example circuit implementation of example data forwarding logic **460** implemented in connection with the main cache store queue **212** of FIGS. 3A-4A. While FIG. 4B is described as example data forwarding logic **460** to be implemented in connection with the main

cache store queue **212**, in other examples disclosed herein, any of the examples described in FIG. **4B** may be implemented in connection with the victim cache store queue **216**.

(148) In the example of FIG. **4B**, the example address line **462** (RD_ADDR) is coupled to the cache controller **220** to obtain an example read instruction from the CPU **102**. In the example of FIG. **4B**, such an example address line **462** obtains the read instruction for the main cache store queue **212**. In this manner, the main cache store queue **212** could forward any data from in-flight write transactions when executing the CPU **102** read instruction. In the main cache store queue **212**, an instruction from the CPU **102** (e.g., a read instruction and/or otherwise transaction, and/or a write instruction and/or otherwise transaction) includes an address, a byte enable (byten) and associated data. The topology of FIG. **4B** includes example address stores **464a** (e.g., the pending store address data store **416** of FIG. **4**), example byten stores **464b** (e.g., byte enable stores), and example data stores **464c** (e.g., any of the latches **402a**, **402b**, **402c**, etc.). In FIG. **4B**, address **addr0** of the address stores **464a** corresponds to the oldest read and/or write instruction in the main cache store queue **212** and address **addr3** of the address stores **464a** corresponds to the newest read and/or write instruction in the main cache store queue **212**.

(149) The incoming address via the address line **462** (e.g., a read instruction) is compared against all the addresses (e.g., **addr0**, **addr1**, **addr2**, and **addr3**) in the address stores **464a** by example compare logic **466**. In examples disclosed herein, the compare logic **466** may compare the address of the read instruction against any number of entries as addresses **addr0**, **addr1**, **addr2**, and **addr3** may be fully or partially equivalent (e.g., the same). In this manner, the example outputs of the compare logic **466** are 4-bit signals. Such 4-bit signals are generated based on the corresponding byte enable signals of the byten stores **464b**. Such 4-bit signals may have any value (e.g., all zeros, all ones, partial zeros and partial ones, etc.).

(150) In example operation, the four addresses of the address stores **464a** may include one or more valid bytes (e.g., bytes that are a logic high). In examples disclosed herein, the main cache store queue **212** is 64 bits (e.g., 8 bytes) and, thus, the main cache store queue **212** may perform a write on any number of bytes, such as, from one to eight.

(151) In response to the compare logic **466** identifying the address of the address stores **464a** that are to be included based on the byte enables in the byten stores **464b**, the example output lines **468** transmit a 8-bit result for each address in the address stores **464a**. In examples disclosed herein, there are four output lines **468** (e.g., **dram_hit_DP0[0][7:0]**, **dram_hit_dp0[1][7:0]**, **dram_hit_dp0[2][7:0]** and **dram_hit_dp0[3][7:0]**). In operation, bits set to a logic high (e.g., 1) in any of the output lines **468** indicate that the corresponding byte of the corresponding address of the address stores **464a** has valid data to be forwarded. For example, if the first output line of the output lines **468** includes **dram_hit_dp0 [0]**, the byte value for the corresponding read instruction may be obtained from any of the addresses in the address store **464a**.

(152) The example of FIG. **4B** includes example select logic **470** and example select lines **472**. In examples disclosed herein, the select lines **472** include eight, 2-byte outputs. Accordingly, there is one select signal of the select signals for each of the corresponding byte enables of the byten stores **464c**. In an example operation, in the event a read instruction indicates to read multiple addresses having multiple enabled bytes, the example select logic **720** selects the most recent data. Such an output of the select logic **470** (e.g., the select lines **472**) control example multiplexers **474a-474h** (multiplexers **474e-h** not shown), respectively. In the example of FIG. **4B**, the multiplexers **474a-h**, include 8 1-byte input terminals. In operation, the multiplexers **474a-h** obtain their respective select line of the select lines **472** which indicate which byte of each of the data stores **464c** to forward. In this manner, the main cache store queue **212** can forward data from different in-flight stores (e.g., any of the data stores **464c**) based on address comparison (e.g., a comparison performed by the comparators **466**), their corresponding byte enables of the byten stores **464b**, and the order in which the main cache store queue **212** was issued a write instruction (e.g., the comparison performed by the select logic **470**).

(153) In examples disclosed herein, any of the address line **462**, the address stores **464a**, the byten stores **464b**, the data stores **464c**, the compare logic **466**, the output lines **468**, the select logic **470**, the select lines **472**, and/or the multiplexers **474a-h** may be implemented in the victim cache store queue **216**. In FIG. **4B**, the output terminals of the multiplexers **474a-h** are coupled to an example cache multiplexers **476**. In the example of FIG. **4B**, cache multiplexer **476** is also coupled to similar multiplexers implemented in this manner in association with the victim cache store queue **216**. In operation, the cache multiplexer **476** obtains a select signal from the cache controller (e.g., the main cache controller **222** or the victim cache controller **224**) that transmitted the read instruction. In this manner, the cache multiplexer **476** facilitates data forwarding to the CPU **102**.

(154) In some examples, the topology of FIG. **4B** may correspond to the example write data forwarding component of the main cache store queue **212**, the write data forwarding component of the victim cache store queue **216**, and the MUX circuits **314**, **316**. In such an example, the cache multiplexer **476** may be implemented by the example MUX circuits **314**, **316**. Additionally, in such an example, any of the address line **462**, the address stores **464a**, the byten stores **464b**, the data stores **464c**, the compare logic **466**, the output lines **468**, the select logic **470**, the select lines **472**, and/or the multiplexers **474a-h** may be implemented by the example write data forwarding component of the main cache store queue **212**, and of the address line **462**, the address stores **464a**, the byten stores **464b**, the data stores **464c**, the compare logic **466**, the output lines **468**, the select logic **470**, the select lines **472**, and/or the multiplexers **474a-h**, as implemented in association with the victim storage **216**, may be implemented by the example write data forwarding component of the victim cache store queue **216**.

(155) In some examples, the topology of FIG. **4B** may correspond to the pending store address data store **418** and example priority multiplexer **418**. For example, the address stores **464a**, the byten stores **464b**, and/or the data stores **464c** may be implemented by the example pending store address data store **416**. In a further example, any of the address line **462**, the compare logic **466**, the output lines **468**, the select logic **470**, the select lines **472**, and/or the multiplexers **474a-h** may be implemented by the example priority multiplexer **418**.

(156) In examples disclosed herein, the topology of FIG. **4B** is utilized for each bank of the main storage **214** and the victim storage **218**. For example, if the main storage **214** has 8 banks, the topology of FIG. **4B** would be replicated 8 times, one for each bank.

(157) Example methods, apparatus, systems, and articles of manufacture to facilitate fully pipelined read-modify-write support in level 1 data cache using store queue and data forwarding are disclosed herein. Further examples and combinations thereof include the following:

(158) Example 1 includes an apparatus comprising a first storage, a second storage, a store queue coupled to the first storage and the second storage, the store queue operable to receive a first memory operation specifying a first set of data, process the first memory operation for storing the first set of data in at least one of the first storage and the second storage, receive a second memory operation, and prior to storing the first set of data in the at least one of the first storage and the second storage, feedback the first set of data for use in the second memory operation.

(159) Example 2 includes the apparatus of example 1, wherein the second memory operation is a read modify-write operation and specifies a second set of data, and the store queue is operable to prior to storing the first set of data in the at least one of the first storage and the second storage, merge the first set of data and the second set of data to produce a third set of data, and provide the third set of data for storing in at least one of the first storage and the second storage.

(160) Example 3 includes the apparatus of example 2, further including receiving a third memory operation indicating to read the third set of data.

(161) Example 4 includes the apparatus of example 1, wherein the second storage is configured as a victim storage of the first storage.

(162) Example 5 includes the apparatus of example 1, wherein the store queue further includes a datastore configured to store a first address value of the first set of data.

- (163) Example 6 includes the apparatus of example 1, wherein the store queue further includes an error code correcting generator to generate error detection code.
- (164) Example 7 includes the apparatus of example 6, wherein the error code correcting generator is to generate the error detection code in response to processing the first memory operation.
- (165) Example 8 includes a method comprising receiving a first memory operation specifying a first set of data, processing the first memory operation for storing the first set of data in at least one of a first storage and a second storage, receiving a second memory operation, and prior to storing the first set of data in the at least one of the first storage and the second storage, transmitting the first set of data for use in the second memory operation.
- (166) Example 9 includes the method of example 8, wherein the second memory operation is a read modify-write operation and specifies a second set of data, the method further including prior to storing the first set of data in the at least one of the first storage and the second storage, merging the first set of data and the second set of data to produce a third set of data, and providing the third set of data for storing in at least one of the first storage and the second storage.
- (167) Example 10 includes the method of example 9, further including receiving a third memory operation indicating to read the third set of data.
- (168) Example 11 includes the method of example 8, wherein the second storage is configured as a victim storage of the first storage.
- (169) Example 12 includes the method of example 8, further including storing a first address value of the first set of data.
- (170) Example 13 includes the method of example 8, further including generating error detection code.
- (171) Example 14 includes the method of example 13, further including generating the error detection code in response to processing the first memory operation.
- (172) Example 15 includes a system comprising a central processing unit configured to transmit a first memory operation and a second memory operation, a first storage coupled to the central processing unit, a second storage coupled to the central processing unit, a store queue coupled to the first storage, the second storage, and the central processing unit, the store queue operable to receive the first memory operation specifying a first set of data, process the first memory operation for storing the first set of data in at least one of the first storage and the second storage, receive the second memory operation, and prior to storing the first set of data in the at least one of the first storage and the second storage, feedback the first set of data for use in the second memory operation.
- (173) Example 16 includes the system of example 15, wherein the second memory operation is a read modify-write operation and specifies a second set of data, and the store queue is operable to prior to storing the first set of data in the at least one of the first storage and the second storage, merge the first set of data and the second set of data to produce a third set of data, and provide the third set of data for storing in at least one of the first storage and the second storage.
- (174) Example 17 includes the system of example 16, further including receiving a third memory operation indicating to read the third set of data.
- (175) Example 18 includes the system of example 15, wherein the second storage is configured as a victim storage of the first storage.
- (176) Example 19 includes the system of example 15, wherein the store queue further includes a datastore configured to store a first address value of the first set of data.
- (177) Example 20 includes the system of example 15, wherein the store queue further includes an error code correcting generator to generate error detection code in response to processing the first memory operation.
- (178) Methods and Apparatus to Reduce Read-Modify-Write Cycles for Non-Aligned Writes
- (179) When issuing a write instruction to a multi-banked storage, such a write instruction may be transmitted with a corresponding read instruction, regardless of the size of the write instruction, in

an attempt to execute a full read-modify-write cycle of such a write instruction. For example, in such an application, a write instruction may be obtained by a CPU indicating to write 128 bits across two 64-bit memory banks, starting at address A0 of the first memory bank. In such an example, though the entirety of both the example memory banks are being written, such an application maintains a read instruction to read the data currently stored in the two example memory banks. However, such an approach is inefficient as twice the processing power (e.g., a write and a read instruction) is used. Additionally, such an approach does not provide any control logic and/or processing circuitry to analyze the write instruction.

(180) In the example illustrated in FIG. 4A, the main storage **214** and/or the victim storage **218** may be multi-banked storages. For example, the main storage **214** may include sixteen memory banks (e.g., sixteen sub-RAMs), each 64 bits wide. In such an example, in the event the cache controller **220** transmits a write instruction to write all 64 bits of a first bank of the main storage **214** (e.g., write a 64-bit word starting with the first address of the first bank), the write instruction can be executed without initiating a read instruction. For example, the bank processing logic **303** may detect that such a write of an entire bank is to be performed and, thus, indicate to the cache controller **220** to initiate the read-modify-write operation, negating to transmit the read instruction.

(181) Similarly, in the event the cache controller **220** transmits a write instruction to the write port **426** to write all 128 bits of a first bank and a second bank of the main storage **214** (e.g., a write instruction indicating to write a 128 bit word starting with the first address of the first bank and ending with the last address of the second bank), the write instruction can be implemented without initiating a read instruction. For example, the bank processing logic **303** may detect that such a write of the entirety of multiple banks is to be performed and, thus, indicate to the cache controller **220** to initiate the read-modify-write operation, negating to transmit the read instruction.

(182) However, in some examples disclosed herein, the cache controller **220** may transmit a write instruction to write 130 bits of data (or any write instruction indicating to write to a subset of the memory banks). Of the 130 bits of data, 64 bits of data may be written to a first bank, 64 bits of data may be written to a second bank, and 2 bits of data may be written to a third bank of the main storage (e.g., a write instruction indicating to write a 130 bit word starting with the first address of the first bank and ending with the second address of the third bank). In such an example, the bank processing logic **303** detects that all addresses of the first bank and the second bank of the main storage **214** are to be written entirely and, thus, indicate to the cache controller to initiate the read-modify-write operations for the first bank and the second bank of the main storage, negating to transmit the read instruction. In such an example, the bank processing logic **303** may detect (e.g., determine) that a subset of the memory banks of the main storage **214** (e.g., the third bank of the memory storage) is to be partially written (e.g., two addresses of the 64 addresses are to be written), and, thus, indicate to the cache controller **220** to initiate a full read-modify-write operation of the third bank of the main storage **214**. In examples disclosed herein, the bank processing logic **303** determines whether to cause a read operation to be performed (e.g., whether to initiate a full read-modify-write operation) in response to the write operation based on whether a number of addresses in the subset of the plurality of memory banks to write satisfies a threshold. In examples disclosed herein, the threshold is not satisfied when the number of addresses in the subset of the plurality of memory banks is greater than 0 and/or less than the number of addresses in the memory bank. In this example, the bank processing logic **303** generates an indication to the CPU **102** to execute the write instruction as a full read-modify-write transaction. In examples disclosed herein, the threshold is satisfied when the number of addresses in the subset of the plurality of memory banks is equal to the number of addresses in the memory bank. In this example, the bank processing logic **303** generates an indication to the CPU **102** to execute the write instruction as a partial read-modify-write transaction (e.g., negating the read). Example description of a read-modify-write operation is described above.

(183) Example methods, apparatus, systems, and articles of manufacture to reduce read-modify-

write cycles for non-aligned writes are disclosed herein. Further examples and combinations thereof include the following:

(184) Example 1 includes an apparatus comprising a memory that includes a plurality of memory banks, an interface configured to be coupled to a central processing unit, the interface to obtain a write operation from the central processing unit, wherein the write operation is to write a subset of the plurality of memory banks, and bank processing logic coupled to the interface and to the memory, the bank processing logic to determine the subset of the plurality of memory banks to write based on the write operation, and determine whether to cause a read operation to be performed in response to the write operation based on whether a number of addresses in the subset of the plurality of memory banks to write satisfies a threshold.

(185) Example 2 includes the apparatus of example 1, wherein the threshold is satisfied when all addresses of the at least one of the number of the plurality of memory banks is included in the write operation.

(186) Example 3 includes the apparatus of example 1, wherein the bank processing logic is to when the number of addresses of the at least one of the number of the plurality of memory banks to write does not satisfy the threshold, generate a second indication to the central processing unit to execute the write operation to the at least one of the number of the plurality of memory banks with issuing the read operation.

(187) Example 4 includes the apparatus of example 1, wherein the bank processing logic determines the number of the plurality of memory banks to write based on the write operation by determining a total number of addresses included in the write operation, and determining the number of the plurality of memory banks based on having an address included in the total number of addresses.

(188) Example 5 includes the apparatus of example 1, wherein the write operation indicates a number of the plurality of memory banks to rewrite.

(189) Example 6 includes the apparatus of example 1, wherein the plurality of memory banks are sixteen memory banks.

(190) Example 7 includes the apparatus of example 1, wherein the memory is a victim storage.

(191) Example 8 includes a system comprising a central processing unit configured to generate a write operation indicating to write to a subset of a plurality of memory banks, a victim storage including the plurality of memory banks, and bank processing logic coupled to the central processing unit and to the victim storage, the bank processing logic to determine the subset of the plurality of memory banks to write based on the write operation, and determine whether to cause a read operation to be performed in response to the write operation based on whether a number of addresses in the subset of the plurality of memory banks to write satisfies a threshold.

(192) Example 9 includes the system of example 8, wherein the threshold is satisfied when all addresses of the at least one of the number of the plurality of memory banks is included in the write operation.

(193) Example 10 includes the system of example 8, wherein the bank processing logic is to when the number of addresses of the at least one of the number of the plurality of memory banks to write does not satisfy the threshold, generate a second indication to the central processing unit to execute the write operation to the at least one of the number of the plurality of memory banks with issuing the read operation.

(194) Example 11 includes the system of example 10, wherein the central processing unit is further to generate a read-modify-write operation to be executed by the at least one of the number of the plurality of memory banks in response to the second indication.

(195) Example 12 includes the system of example 8, wherein the bank processing logic determines the number of the plurality of memory banks to write based on the write operation by determining a total number of addresses included in the write operation, and determining the number of the plurality of memory banks based on having an address included in the total number of addresses.

(196) Example 13 includes a method comprising determining a subset of a plurality of memory banks of a memory to write based on a write operation, and determine whether to cause a read operation to be performed in response to the write operation based on whether a number of addresses in the subset of the plurality of memory banks to write satisfies a threshold.

(197) Example 14 includes the method of example 13, further including transmitting an indication to a central processing unit in response to the threshold being satisfied.

(198) Example 15 includes the method of example 13, wherein the threshold is satisfied when all addresses of the at least one of the number of the plurality of memory banks is included in the write operation.

(199) Example 16 includes the method of example 13, further including when the number of addresses of the at least one of the number of the plurality of memory banks to write does not satisfy the threshold, generating a second indication to execute the write operation to the at least one of the number of the plurality of memory banks with issuing the read operation.

(200) Example 17 includes the method of example 13, wherein determining the number of the plurality of memory banks to write is performed by determining a total number of addresses included in the write operation, and determining the number of the plurality of memory banks based on having an address included in the total number of addresses.

(201) Example 18 includes the method of example 13, wherein the write operation indicates a number of the plurality of memory banks to rewrite.

(202) Example 19 includes the method of example 14, further including obtaining the write operation from the central processing unit, the write operation indicating the number of the plurality of memory banks to write.

(203) Example 20 includes the method of example 14, wherein the memory is a victim storage.

(204) Aggressive Write Merging to Reduce Bank Pressure

(205) In the example of FIG. 4A, the example main cache store queue **212** stores a number of write operations at different sections of the main cache store queue **212** (e.g., at the example latches **402a-e**). For example, when the CPU **102** transmits three separate write operations in a row, the first write operation that the CPU **102** provided is stored at the first latch **402b** and moved to the second latch **402c** when the second operation is received at the first latch **402b**. Accordingly, after the three write operations are received, the first latch **402b** will store and/or output the last write operation with respect to time (e.g., which is last to be stored in the main storage **214**), the second latch **402c** will have the second write operation (e.g., which is second to be stored in the main storage **214**), and the third latch **402d** will have the first write operation (e.g., which was the first to be stored in the example main storage **214**). Whenever data in the main cache store queue **212** is next to be stored in the main storage **214**, the example arbitration manager **414** reserves a cycle for the data to be written into the example main storage **214**. Accordingly, during the reserved cycle, the main storage **214** may not be available to perform read operations. However, if the data operations stored in two or more of the latches **402b**, **402c**, **402d** correspond to the same memory address, the data can be merged in order to write the data into the memory address of the main storage **214** once, instead of two or three times. For example, if the write operation stored in the latch **402d** corresponds to writing a byte of the memory address and the write operation stored in the latch **402c** corresponds to writing the same byte to the memory address, the second write will overwrite the first write.

(206) Instead of reserving a cycle to write the first byte into the main storage **214**, which will be overwritten in a subsequent cycle, the main cache store queue **212** merges the two writes into one write, so that only one cycle is used to write the second transaction (e.g., to avoid reserving a cycle for the first write). Such an aggressive merge reduces the number of cycles reserved for write operations. In this manner, the main storage **214** will have extra cycles to perform read operations, thereby decreasing the latency of the overall systems.

(207) To enable an aggressive merge of write operations to the same memory address, the output of

the example latches **402b-402d** are coupled to the example merging circuits **403a-403c**. For example, the output of the third latch **402d** may be coupled to the merging circuit **403a**, the output of the second latch **402c** may be coupled to the merging circuit **403b**, and the output of the first latch **402b** may couple to the merging circuit **403c**. The output of the merging circuit **403a** may additionally be coupled to the output of the second latch **402c** and the merging circuit **403b**, the merging circuit **403b** may be coupled to the merging circuit **403c**, and the merging circuit **403c** may be coupled to the input of the first latch **402b**. Additionally or alternatively, there may be additional comparator circuits that compare outputs of a different set of latches (e.g., a comparator circuit to compare the output of latch **402b** with the output of latch **402d**). The example merging circuits **403a-c** include example comparator(s) **420** and example switches **422**.

(208) The comparator(s) **420** compare the memory address locations for each write operation that is stored in the respective latches **402b-402d** to determine whether any of the write operations in the example store queue correspond to the same memory address. The example comparator **420** may be one comparator to compare all the write operations of the latches **402b-402d** or may be separate comparators **420**, to compare two of the latches **402b-d** (e.g., a first comparator to the memory address of latch **402b** to the memory address of latch **402c**, a second comparator to the memory address of **402b** to the memory address of latch **402d**, etc.). The comparator(s) **420** output the results of the comparisons (e.g., with one or more signals corresponding to the one or more comparisons) to the example switch(es) **422** and/or the arbitration manager **414**. If the example arbitration manager **414** receives a signal indicative of a match, the arbitration manager **414** will not reserve the cycle for a first write operation while the first write operation is merged with a second write operation to the same memory location (e.g., to free up cycles for other cache operations).

(209) The example switch(es) **422** reroute the write operations in the example latches **402b-402d** based on the comparison. For example, if the memory address of the write operation stored in the example latch **402d** is the same as the memory address stored in the latch **402c**, the example switch(es) **422** enable and/or disable to reroute the output of the latch **402d** to latch **402c**, instead of routing to the example arbitration manager **414**. In this manner, the two write operations are combined and written into the main storage **214** in a subsequent cycle as a single write operation instead of two write operations. The switch(es) **422** may be electrical switches, transistors (e.g., MOSFETS), demultiplexers, and/or any other component that can reroute a signal in a circuit.

(210) When a MUX of one of the merging circuit **403a-c** receives a rerouted write operation, the MUX performs a merging protocol for the one or more rerouted write operations that prioritizes the newest write operation. For example, if the comparator(s) **420** determines that the write operation stored in the example latch **402c** corresponds to the same memory address as the write operation stored in the example latch **402d**, the MUX of the comparator(s) **420** merge the data (as further described in conjunction with FIG. 4D) and the switches(es) **422** reroute the merged write operation stored in the example latch **402d** to the latch **402c**. In such an example, if the write operation in the example latch **402c** corresponds to writing data to byte0 and byte2 of the memory address and the write operation from the latch **404d** corresponds to writing data to byte0 and byte1, the example merging circuit **403a** merges the two write operations to keep the writing data stored in latch **402c** (e.g., the write to byte0 and byte2) and include the write data from latch **402d** that doesn't overlap (e.g., byte2). In this example, the merging circuit **403a** discards the write data of byte 0 from the latch **404d** as part of the merging operation because the data to be written at byte 0 from the latch **404d** will be overwritten by the write instructions of the latch **402c**. Accordingly, in the above example, the merged data corresponds to the write data for byte0 from latch **402c**, the write data for byte1 from latch **404d**, and the write data for byte2 from the latch **402c**. When a subsequent write operation is added to the main cache store queue **212**, the merged write data from the latch **402c** may be manipulated (e.g., via one of the example blocks **404**, **406**, **408**) and/or pushed to the next latch **402d** to be stored in the example main storage **214** during a subsequent

cycle. An example hardware implementation of the merging protocol is further described below in conjunction with FIG. 4D.

(211) FIG. 4D illustrates a hardware implementation of the merging circuit **402c** of FIG. 4A and/or **1102c** of FIG. 11A (e.g., to merge data from the latch **402d** to the latch **402c**). Similar hardware setups can be implemented to merge data between any two latches. The example of FIG. 4D includes the example latches (e.g., stores) **402b-402d** and hardware components for the comparator **420** and the switch **422** of FIG. 4A. The example comparator circuit **420** includes comparators and/or logic gates **480a-480f** and the switch circuit **422** includes an OR gate **482** and a MUX **484**. Although described in conjunction with FIG. 4A, if described in conjunction with the victim cache store queue **216** of FIG. 11A, the latches **1102a-d**, example comparator **1120** and example switch **322** of FIG. 11A could be used.

(212) In the example of FIG. 4D, the example latch **402d** outputs the stored data to the example storage (e.g., the main storage **214** or the victim storage **218** via the arbitration manager **414**, **1114**), which locks its bank for a first cycle. FIG. 4D illustrates the write-merge locking old data in the bank of the storage when the old data is preceded by another store/latch including new write data to the same address. Although the example of FIG. 4C illustrates a merge between the example latch **402d** and the example latch **402c**, a similar structure may be used to merge data between any of the latches **402a-402d**. For example, if the address of data at the latch **402d** matches with the address of the data at the latch **402b**, then the latch **402** can merge its data with the data at latch **402b**. Additionally, the data at three or more latches may be merged into a single latch if the data at the three or more latches correspond to the same address. As described above, if the data at a particular latch is merged with data and a prior latch, the data at the particular latch is invalid (e.g., by setting a bit to a value corresponding to invalidated) or discarded so that the arbitration is not performed for that data to be locked in a bank in the storage.

(213) In operation, the example comparator **480b** compares the address of the data at latch **402d** is compared to the address of the data at latch **403c**. The output of the comparator **480b** is valid only if both the latches **402c** and **402d** include a bit that is set to valid. Accordingly, the example comparator **480a** compares the valid bits from the latches **402c**, **402d** to ensure that the bits are both valid. In this manner, the example comparator **480d** only outputs a value (e.g., '1') when the address at latch **402c** and the address at latch **402d** are the same and the valid bits of both latches **402c**, **402d** are valid. Otherwise, the example comparator **480d** outputs a second value ('0'). The output of the comparator **480d** is input to the select line of the MUX **484** of the switch **422**. Thus, if the addresses of the latches **402c**, **402d** are the same and value, the comparator **480d** will control the MUX **484** of the switch **422** to output merged data (e.g., from the OR gate **482** of the switch **422**).

(214) The example ADD gate **480c** performs a logic AND function with the byten of latch **402d** and not of the byten of latch **402c** (e.g., using a NOT gate). The output of the example ADD gate **480c** ensures that if there is a byte being written by both the latch **402c** and **402d**, then the latch **402c** is held and the latch **402d** is ignored because the latch **402c** has the most recent store issued by the CPU **102**. Accordingly, the example comparator **480f** only outputs the most recent stored issued by the CPU **102** for an address in the store queue. The output of the comparator **480f** corresponds to `dram_mrg_data_sel[7:0]`, which is the signal used to select final merged data between the logic gate **402d** and logic gate **402c** on a byte by byte basis. Similar merge logic exists on merging the byten of the latch **402c** and the latch **402d**. The signal marked as (1) is used to selected between `store_byten` or (`store3_byten` merged with `store_4_byten`), as described above.

(215) Example methods, apparatus, systems, and articles of manufacture to facilitate aggressive write merging to reduce bank pressure are disclosed herein. Further examples and combinations thereof include the following: Example 1 includes an apparatus comprising a first cache storage, a second cache storage, a store queue coupled to at least one of the first cache storage and the second cache storage and operable to receive a first memory operation, process the first memory operation

for storing the first set of data in at least one of the first cache storage and the second cache storage, receive a second memory operation, and prior to storing the first set of data in the at least one of the first cache storage and the second cache storage, merge the first memory operation and the second memory operation.

(216) Example 2 includes the apparatus of example 1, wherein the first memory operation specifies a first set of data, the second memory operation specifies a second set of data, and the store queue is operable to prior to storing the first set of data in the at least one of the first cache storage and the second cache storage, merge the first set of data and the second set of data to produce a third set of data, and provide the third set of data for storing in at least one of the first cache storage and the second cache storage.

(217) Example 3 includes the apparatus of example 2, wherein the store queue is to store the third set of data in the at least one of the first cache storage or the second cache storage in one cycle.

(218) Example 4 includes the apparatus of example 1, wherein the store queue is operable to merge the first memory operation and the second memory operation by cancelling a part of the first memory operation.

(219) Example 5 includes the apparatus of example 4, wherein the part of the first memory operation are bytes that the second memory operation is to write to.

(220) Example 6 includes the apparatus of example 4, wherein the part is a first part, the store queue operable to merge the first memory operation and the second memory operation by maintaining a second part of the first memory operation.

(221) Example 7 includes the apparatus of example 6, wherein the second part of the first memory operation are bytes that the second memory operation is not to write to.

(222) Example 8 includes the apparatus of example 1, wherein the first cache storage is a main cache storage and the second cache storage is a victim cache storage.

(223) Example 9 includes a system comprising a central processing unit coupled in parallel to a first cache storage and a second cache storage, a store queue coupled to at least one of the first cache storage and the second cache storage and operable to process a first memory operation from the central processing unit, the first memory operation for storing the first set of data in at least one of the first cache storage and the second cache storage, prior to storing the first set of data in the at least one of the first cache storage and the second cache storage, merge the first memory operation and a second memory operation corresponding to a same memory address.

(224) Example 10 includes the system of example 9, wherein the first memory operation specifies a first set of data, the second memory operation specifies a second set of data, and the store queue is operable to prior to storing the first set of data in the at least one of the first cache storage and the second cache storage, merge the first set of data and the second set of data to produce a third set of data, and provide the third set of data for storing in at least one of the first cache storage and the second cache storage.

(225) Example 11 includes the apparatus of example 10, further including a store queue to store the third set of data in the at least one of the first cache storage or the second cache storage in one cycle.

(226) Example 12 includes the system of example 9, wherein the store queue is operable to merge the first memory operation and the second memory operation by cancelling a part of the first memory operation.

(227) Example 13 includes the system of example 12, wherein the part of the first memory operation are bytes that the second memory operation is to write to.

(228) Example 14 includes the system of example 12, wherein the part is a first part, the store queue operable to merge the first memory operation and the second memory operation by maintaining a second part of the first memory operation.

(229) Example 15 includes the system of example 14, wherein the second part of the first memory operation are bytes that the second memory operation is not to write to.

(230) Example 16 includes the system of example 9, wherein the first cache storage is a main cache storage and the second cache storage is a victim cache storage.

(231) Example 17 includes a method comprising obtaining a first memory operation, processing, by executing an instruction with a processor, the first memory operation for storing a first set of data in at least one of a first cache storage and a second cache storage, obtaining a second memory operation, and prior to storing the first set of data in the at least one of the first cache storage and the second cache storage, merging, by executing an instruction with the processor, the first memory operation and the second memory operation.

(232) Example 18 includes the method of example 17, wherein the first memory operation corresponds to a first set of data and the second memory operation corresponds to a second set of data, further including prior to storing the first set of data in the at least one of the first cache storage and the second cache storage, merging the first set of data and the second set of data to produce a third set of data, and providing the third set of data for storing in at least one of the first cache storage and the second cache storage.

(233) Example 19 includes the method of example 18, further including storing the third set of data in the at least one of the first cache storage or the second cache storage in one cycle.

(234) Example 20 includes the method of example 17, further including merging the first memory operation and the second memory operation by cancelling a part of the first memory operation.

(235) Optimized Atomic Operations and Histogram Operations Support in L1 Cache

(236) Atomic operations are further examples of multi-part memory operations. For example, an atomic compare and swap operation manipulates a value stored at a memory location based on the results of a comparison of the existing value stored at the memory location. For example, the CPU **102** may want to replace the data stored in the L1 cache **110** with a new value if the existing value stored in the L1 cache **110** matches a specific value.

(237) In some example systems, when a CPU wanted to perform an atomic operation, the CPU sent a read operation to a memory address, performed the manipulation on the read data, and then executed a write operation to the same memory address to store the manipulated data. Additionally, in such example systems, the L1 cache paused, rejected, blocked, and/or halted any transactions from other devices (e.g., other cores of the CPU, higher level cache, the extended memory, etc.) until the atomic operation was complete (e.g., to avoid manipulation of the memory address corresponding to the atomic operation during the atomic operation). Accordingly, such example techniques required lots of effort on behalf of the CPU and lots of reserved cycles that increase latency.

(238) In the example of FIG. 4A, the example main cache store queue **212** handles atomic operations in conjunction with the read modify write structure. In this manner, the example CPU **102** can send a single atomic operation to the L1 cache **110**, and the main cache store queue **212** handles the atomic data manipulation and writing operation. Accordingly, the CPU **102** utilizes a single cycle to execute an atomic operation and can use the other cycles (e.g., used in some example atomic protocols) to perform other functions, thereby reducing the latency of the overall computing system **100**.

(239) In operation, when the CPU **102** transmits an atomic operation and/or an atomic compare and swap operation to increment and/or swap the data at a memory address by a value of 1, for example, the atomic instruction is received by the latch **402a** and the tag RAM **208** verifies whether the memory address is stored in the example main storage **214**. If the memory address is stored in the example main storage **214**, the tag RAM **208** instructs the example main storage **214** to output the data at the memory address while the atomic instructions are passed to the example latch **402b**. While the main storage **214** outputs the data to the latch **322a**, the example latch **402b** outputs the atomic operation to the latch **402c**. At this point the ECC logic **310** performs error detection and/or correction protocol as described above, and the data from the memory address location is forwarded to the example arithmetic component **404** (e.g., for atomic operations) or the

atomic compare component **406** (e.g., for the atomic compare and swap operations). The arithmetic component **404** obtains the atomic operation (e.g., including data identifying how to manipulate the data) and/or the atomic compare and swap **406** obtains the atomic compare and swap operation (e.g., including a key and data to be written if the key matches read data) from the latch **402c** and obtains the data from the corresponding memory address from the output of the ECC logic **310**. The arithmetic component **404** performs the manipulation to the data (e.g., increment the data by 1) and/or the atomic compare component **406** may perform the swap (replaces the data if the read data matches a key, etc.) and outputs the incremented and/or swapped-in data for the corresponding memory address (e.g., the atomic result) to the example latch **402d** via the example MUX **410**. (e.g., which is enabled via the cache controller **220**). The latch **402d** outputs the new data corresponding to the memory address to the ECC generator **412** to generate the ECC bit and the arbitration manager **414** writes the new data (e.g., the atomic result and/or atomic compare and swap result) to the memory address in conjunction with the ECC bit in the example main storage **214**. Additionally or alternatively, the corrected value out of the EDD logic **412** is returned to the CPU **102**. Thus, the atomic operation is performed with only one instruction from the CPU **102**.

(240) The atomic compare component **406** and/or the arithmetic component **404** have several inputs. For example, the atomic component **406** receives (e.g., obtains) the type of atomic operation to perform (e.g. atomic compare and swap, or atomic swap), the new data to swap in, the ECC corrected data read out the of the cache **310**, and the size of the size of the data to be manipulated during the atomic operation (e.g., 32-bit or 64-bit). In the example circuit implementation **450** of FIG. **4B**, the atomic compare component **406** receives an atomic compare and swap operation and the arithmetic component **404** receives an atomic operation. The atomic compare component **406** compares the comparison value (e.g., a key) provided by the CPU **102** against the ECC data **310**. On a match, the new data is swapped in place of the old data (e.g. ECC data **310**) and output to the MUX **410**. The size of the new data swapped-in is determined by `cas_acc_sz` input (e.g. 32-bit or 64-bit). In the example circuit implementation **450** of FIG. **4C**, the atomic compare component **406** may also receive an atomic swap operation. The atomic compare component **406** will swap-in the new data replacing the ECC data **310** regardless of the comparison result and output the new value to the mux **410** and the old data from the address is read from the main storage **214** and is provided back to the CPU **102**. The size of the new data swapped-in is determined by `cas_acc_sz` input (e.g. 32-bit or 64-bit). In the example circuit implementation **450** of FIG. **4C**, the arithmetic component **404** may also receive an atomic operation. The arithmetic component **404** will manipulate the ECC data **310** and store the manipulated data in the main storage element **214**. The size of the new data swapped-in is determined by `cas_acc_sz` input (e.g. 32-bit or 64-bit). Although the example circuit **450** of FIG. **4C** is described above in conjunction with the main store queue, the circuit **450** can be described in conjunction with a victim store queue as further described below in conjunction with FIG. **11A**.

(241) FIG. **4C** illustrates an example circuit diagram of parts of the main cache store queue **212** of FIG. **4A** and/or parts of the victim cache store queue **216** of FIG. **11A**. FIG. **4C** illustrates a detailed circuit diagram of the arithmetic unit **404**, **1104**.

(242) The arithmetic unit **404** may be used for other types of memory transactions such as histogram operations. A histogram operation retrieves a value stored in memory that may represent a bin of a histogram, the histogram operation then modifies the values before storing it back to the same memory address or an alternative address. In an example, a first data set contains the values [0, 0, 2, 0, 0, 3], and a second data set contains bins representing the number of occurrences of respective values within the first data set. In some example systems that lack an arithmetic unit **404**, to perform a histogram function, the CPU reads each in the first data set and increments the second data set for each value. For example, to determine how many 0s are stored in a 10 byte data set, the CPU may perform 10 reads. Then to determine how many is are in the same data set, the CPU will perform an additional 10 reads. Thus, to perform a histogram operation, such systems

may need to perform (N)(M) reads, where N is the size of the section of memory (e.g., 10 bytes) being read and M is the number of values that could be store in each byte Like an atomic operation, the L1 SRAM may have to block, pause, halt, discard, etc. all other read and/or write operations until the histogram operation is complete. In contrast, the arithmetic unit **404** may be used to perform the same operation with a single transaction from the CPU.

(243) The arithmetic unit **404**, **1104** has several inputs including the type of histogram (e.g., weighted or increment, signed or not signed), histogram size (e.g., byte, halfword, or word), histogram weight (e.g., signed weights to be added per selected bins for weighted histograms, weight equal to one for normal histogram), histogram sign (e.g., existing big signed or not), and histogram saturate enable.

(244) Component **451** of FIG. **4C** selects a bin read out of the ECC component **310** for the bank illustrated in FIG. **4C**. Component **452** selects the weight to be added to the bin from the vector of weights provided by the CPU **102**. Cnt_value is the sum of the bin value from component **451** and the weight provided by the CPU **102**. Component **453**, component **454** and component **458** are used as part of the saturation circuit. Component **453** receives the histogram size (byte, halfword, or word) and the count value (the sum of the outputs of components **451**, **452**) and determines if a signed bin will saturate. Component **454** receives the histogram size (byte, halfword, or word) and the count value (the sum of the outputs of components **451**, **452**) and determines if a unsigned bin will saturate. The outputs of components **453** and **454** are sent to a MUX which selects whether to use the signed saturate value or the unsigned saturate values based on the kind of bin selected (e.g. signed or unsigned). Component **458** receives the output of the MUX, the histogram saturate enable, the type of histogram (e.g. signed or not signed), the sign bit of the bin, and the sign bit of the weight and outputs an updated cnt_value adjusted for the selected saturation kind (e.g., based on the above Table 1). Components **455**, **456**, **457** are used to position the resultant cnt_value and byte enables in the correct bit positions for writing the new data bank to the correct bin.

(245) In the example of FIG. **4A**, the CPU **102** instructs the main storage **214** to perform the histogram operation. Thereby changing the number of cycles that the CPU **102** has to reserve for the operation from (N)(M) to 1. Additionally, because the atomic operation protocol is already implemented in the store queue, the histogram operation can be performed using the arithmetic component **404** by performing N reads for the N size of the memory and incrementing a count for each value in the example main cache store queue **212**, thereby reducing the number or read operation from (N)(M) operations to N operations.

(246) In operation, when the CPU **102** transmits a histogram operation corresponding to a section (e.g., a SRAM line) of the main storage **214**, the histogram operation is stored in the example latch **402a** while the tag RAM **208** verifies whether the memory address corresponding to the histogram operation is available in the main storage **214**. The example cache controller **220** facilitates the read operation for each byte of the section identified in the histogram operation (e.g., where histogram bins are accessed in parallel by reading up to 128 Bytes at the same time). If available, the tag RAM **208** instructs the main storage **214** to output the data at a first byte of the section of the main storage **214** while the histogram operation is output by the example latch **402a** to the example latch **402b**. When the example main storage **214** outputs the data that has been read from the memory address to the example latch **322a**, the latch **402b** outputs the histogram operation to the example latch **402c**. After the ECC logic **310** performs the error detection and correction functionality, the data read at the byte is sent to the example arithmetic component **404**.

(247) After receiving the read value from the ECC logic **310** and the histogram instructions from the latch **402c**, the arithmetic component **404** initiates data representative of the histogram. For example, the arithmetic component **404** may initiate a vector (e.g., representing a histogram) with an initial value (e.g., zero) for each possible value that could be stored in the bytes of the main storage. The arithmetic component **404** increments the value of the vector based on output by the ECC logic **310** (e.g., the read byte). For example, if the read value of the byte is 0, the arithmetic

component **404** increments the value corresponding to 0 in the vector.

(248) Because each vector is incremented in parallel, the resulting vector corresponds to a histogram of the values that were read in the corresponding sections of SRAM in parallel. Because a value of the histogram is incremented for each bit, the resulting vector is a histogram of the values stored in the section of memory identified in the histogram operation from the CPU **102**. In some examples, the arithmetic component **404** may increment in parallel by some weighted value (e.g., 1.5). When complete, the example histogram is input to the example MUX **418** (e.g., controlled by the example pending store address table **416**) to be input to the MUX **314** via the example latch **402e**. Once the operation is complete, the example cache controller **220** controls the MUX **314** to output the final histogram vector to the example CPU interface **202** via the example latch **322b**, thereby ending the histogram operation.

(249) In some examples, the L1 cache **110** supports functionality where a histogram bin can saturate after the histogram bin includes more than a threshold limit of the bin size (e.g., a byte, a halfword, a word, etc.). The below Table 1 illustrates an example of saturation values. Using this functionality, the histogram bin values will not roll over once they reach the maximum value.

(250) TABLE-US-00001

TABLE 1 Histogram Saturation Values		Saturate	Saturate	Unsaturate	
Unsaturate	Min	Max	Min	Max	
Byte	Signed	0x80	0x7F	0x80	0x7F
Byte	Unsigned	0x0	0xFF	0x0	0xFF
HalfWord	Signed	0x8000	0x7FFF	0x8000	0x7FFF
HalfWord	Unsigned	0x0	0xFFFF	0x0	0xFFFF
Word	Signed	0x8000_0000	0x7FFF_FFFF	0x8000_0000	0x7FFF_FFFF

(251) Example methods, apparatus, systems, and articles of manufacture to facilitate optimized atomic and histogram operations are disclosed herein. Further examples and combinations thereof include the following: Example 1 includes a system comprising a cache storage coupled to an arithmetic component, and a cache controller coupled to the cache storage, wherein the cache controller is operable to receive a memory operation that specifies a set of data, retrieve the set of data from the cache storage, utilize the arithmetic component to determine a set of counts of respective values in the set of data, generate a vector representing the set of counts, and provide the vector.

(252) Example 2 includes the system of example 1, wherein the cache controller is operable to provide the vector for storing in the cache storage.

(253) Example 3 includes the system of example 1, wherein the cache controller is operable to provide the vector to a processor.

(254) Example 4 includes the system of example 1, further including a store queue coupled to the cache controller, the store queue including the arithmetic component.

(255) Example 5 includes the system of example 1, wherein the cache storage is at least one of a main cache storage or a victim cache storage.

(256) Example 6 includes the system of example 1, wherein the arithmetic component is to obtain (a) the set of data from the cache storage via an error detection and correction circuit and (b) the memory operation from a central processing unit via a latch.

(257) Example 7 includes the system of example 1, wherein the cache controller is operable to provide the vector to a central processing unit based on a single instruction from the central processing unit at a single cycle.

(258) Example 8 includes a system comprising a cache storage, and a cache controller coupled to the cache storage and an arithmetic component, wherein the cache controller is operable to receive a memory operation specifying a first set of data and an arithmetic operation, retrieve the first set of data from the cache storage, utilize the arithmetic component to perform the arithmetic operation on the first set of data to produce a second set of data, and provide the second set of data.

(259) Example 9 includes the system of example 8, wherein the cache controller is operable to provide the second set of data for storing in the cache storage.

(260) Example 10 includes the system of example 8, wherein the cache controller is operable to provide the second set of data to a processor.

- (261) Example 11 includes the system of example 8, further including a store queue coupled to the cache controller, the store queue including the arithmetic component.
- (262) Example 12 includes the system of example 8, wherein the cache storage is at least one of a main cache storage or a victim cache storage.
- (263) Example 13 includes the system of example 8, wherein the arithmetic component is to obtain (a) the first set of data from the cache storage via an error detection and correction circuit and (b) the memory operation from a central processing unit via a latch.
- (264) Example 14 includes the system of example 8, wherein the cache controller is operable to provide the second set of data to a central processing unit based on a single instruction from the central processing unit at a single cycle.
- (265) Example 15 includes a method comprising obtaining a memory operation that specifies a set of data, obtaining the set of data from a cache storage, determining a set of counts of respective values in the set of data, generating a vector representing the set of counts, and providing the vector.
- (266) Example 16 includes the method of example 15, wherein the vector is provided to the cache storage.
- (267) Example 17 includes the method of example 15, wherein the vector is provided to a processor.
- (268) Example 18 includes the method of example 15, wherein the cache storage is at least one of a main cache storage or a victim cache storage.
- (269) Example 19 includes the method of example 15, further including obtaining (a) the set of data from the cache storage via an error detection and correction circuit and (b) the memory operation from a central processing unit via a latch.
- (270) Example 20 includes the method of example 15, wherein the providing the vector to a central processing unit is based on a single instruction from the central processing unit at a single cycle.
- (271) Atomic Compare and Swap Support in L1 Cache for Coherent System
- (272) If data from an address is stored in the L1 cache **110**, data for the same address may also be cached in other caches (e.g., L2 cache **112**, and/or L3 cache **114**). Accordingly, the memory address needs to be tracked and/or monitored to ensure that any changes to the data in one cache is identified, tracked, and/or reflected to the other caches (e.g., L2 cache **112** and L3 cache **114**). As described above in conjunction with FIGS. **3A-3D**, the example MESI RAM **300** tracks the state of the data stored in the main storage **214** to be able to avoid issues with mismatched data in different caches that correspond to the same memory address. For example, if the CPU **102** transmits a read operation, the example MESI RAM **300** changes the state of the memory address to shared, because the data in the memory address will not be manipulated. If the CPU **102** transmits a write operation, the example MESI RAM **300** changes the state of the memory address to exclusive, because the data in the memory address will be manipulated and the main storage **214** needs write permission for the address. After the data in the memory address is written to the main storage **214**, the MESI RAM **300** updates the state of the memory address to modified (e.g., indicating that the memory address has been modified).
- (273) As described above, during atomic and/or histogram protocols, the data from a memory address is read from the main storage **214** and provided to the main cache store queue **212** to be updated (e.g., incremented) and written back into the main storage **214**. However, if the MESI RAM **300** has identified the state of the corresponding memory address as in shared state, the write operation of the atomic protocol may cause problems with other level caches (e.g., because the write will cause a mismatch of data in different caches).
- (274) To avoid such data mismatches during atomic operations (e.g., including atomic compare and swap), the example cache controller **220** marks cache hits that correspond to a shared state as a cache miss. In this manner, the cache controller **220** can instruct the L2 interface **228** to send the cache miss to the higher level cache with an exclusive state request. In this manner, the higher level

cache can grant the exclusive state to the L1 cache **110** and the L1 cache **110** can perform the read and write operation as part of the atomic operation in response to receiving the granted exclusive state.

(275) After the cache miss with exclusive request is sent to the higher level cache and the higher level cache responds with the write data and the exclusive response, the example atomic operation logic **406** will instruct the MESI RAM **300** to tag the data as modified. The received data from the L2 cache **112** is transmitted into the main cache store queue **212** to be stored in the main storage **214**. Because the operation was an atomic operation (e.g., a regular atomic operation or an atomic compare and swap) or a histogram protocol, the data from the higher level cache is manipulated by the example arithmetic component **404** and/or the example atomic compare component **406** for the manipulation and stored in the example main storage **215** via the example ECC generator **412** and the example arbitration manager **414**.

(276) Example methods, apparatus, systems, and articles of manufacture to facilitate an atomic compare and swap are disclosed herein. Further examples and combinations thereof include the following: Example 1 includes a system comprising a cache storage, a cache controller coupled to the cache storage wherein the cache controller is operable to receive a memory operation that specifies a key, a memory address, and a first set of data, retrieve a second set of data corresponding to the memory address, compare the second set of data to the key, based on the second set of data corresponding to the key, cause the first set of data to be stored at the memory address, and based on the second set of data not corresponding to the key, complete the memory operation without causing the first set of data to be stored at the memory address.

(277) Example 2 includes the system of example 1, further including a storage queue for the cache storage, the storage queue including a latch to interface with a tag random access memory (RAM) to retrieve the second set of data from the cache storage, an atomic compare component, the cache controller to cause the atomic compare component to compare the second set of data to the key, and if the key does not match the second set of data, discard the first set of data, and an arbitration manager to, if the key matches the second set of data, store the first set of data at the memory address in the cache storage.

(278) Example 3 includes the system of example 1, wherein the cache storage is at least one of a main cache storage or a victim cache storage.

(279) Example 4 includes the system of example 2, wherein the latch is a first latch, the atomic compare component to obtain (a) the second set of data from the cache storage via an error detection and correction circuit and (b) the atomic compare and swap command from the first latch via a second latch.

(280) Example 5 includes the system of example 1, wherein the cache controller is operable to complete the memory operation based on a single instruction from a central processing unit at a single cycle.

(281) Example 6 includes the system of example 1, further including a modified, exclusive, shared, invalid (MESI) component to determine a state of the memory address of the memory operation, and the cache controller operable to, if the memory address of the memory operation corresponds to a shared state, provide miss instructions to another cache, the miss instructions including an exclusive state request.

(282) Example 7 includes the system of example 6, wherein the cache controller is to cause an arbitration manager to, if the first data matches the key, store the first set of data at the memory address in the cache storage after the exclusive state request has been granted from the other cache.

(283) Example 8 includes the system of example 6, wherein the cache controller causes a tag RAM to retrieve the second set of data from the cache storage from the memory address after the exclusive state request has been granted from the other cache.

(284) Example 9 includes a storage queue for a cache storage, the storage queue comprising a latch to retrieve a memory operation that specifies a key, a memory address, and a first set of data, and

retrieve a second set of data corresponding to the memory address, and an atomic compare component to compare the second set of data to the key, based on the second set of data corresponding to the key, cause the first set of data to be stored at the memory address, and based on the second set of data not corresponding to the key, complete the memory operation without causing the first set of data to be stored at the memory address.

(285) Example 10 includes the storage queue of example 9, wherein the cache storage is at least one of a main cache storage or a victim cache storage.

(286) Example 11 includes the storage queue of example 9, wherein the latch is a first latch, the atomic compare component to obtain (a) the second set of data from the cache storage via an error detection and correction circuit and (b) the atomic compare and swap command from the first latch via a second latch.

(287) Example 12 includes the storage queue of example 9, wherein, if the first data matches the key, the storage queue performs the atomic compare and swap command based on a single instruction from a central processing unit at a single cycle.

(288) Example 13 includes the storage queue of example 9, further including a modified, exclusive, shared, invalid (MESI) component to determine a state of the memory address of the memory operation, and an interface to, if the memory address of the memory operation corresponds to a shared state, provide miss instructions to another cache, the miss instructions including an exclusive state request.

(289) Example 14 includes the storage queue of example 13, further including an arbitration manager to, if the first data matches the key, store the first set of data at the memory address after the exclusive state request has been granted from the other cache.

(290) Example 15 includes the storage queue of example 13, wherein the latch is to interface with a tag RAM to read the second set of data after the exclusive state request has been granted from the other cache.

(291) Example 16 includes a method comprising receiving a memory operation that specifies a key, a memory address, and a first set of data, retrieving a second set of data corresponding to the memory address, comparing the second set of data to the key, based on the second set of data corresponding to the key, causing the first set of data to be stored at the memory address, and based on the second set of data not corresponding to the key, completing the memory operation without causing the first set of data to be stored at the memory address.

(292) Example 17 includes the method of example 16, further including interfacing with a tag random access memory (RAM) to retrieve the second set of data from a cache storage, comparing the second set of data to the key, and if the key does not match the second set of data, discarding the first set of data, and if the key matches the second set of data, storing the first set of data at the memory address in the cache storage.

(293) Example 18 includes the method of example 17, further including (a) the second set of data via an error detection and correction circuit and (b) the memory operation from a central processing unit.

(294) Example 19 includes the method of example 16, wherein the completing of the memory operation is based on a single instruction from a central processing unit at a single cycle.

(295) Example 20 includes the method of example 16, further including determining a state of the memory address of the memory operation, and if the memory address of the memory operation corresponds to a shared state, providing miss instructions to another cache, the miss instructions including an exclusive state request.

(296) Methods and Apparatus for Inflight Data Forwarding and Invalidation of Pending Writes in Store Queue

(297) In examples disclosed herein, when the CPU **102** issues a read operation, the tag RAM accesses **204**, **206** determine if the read operation hits or misses the main storage element **214** and/or the victim storage element **218**. In examples where the tag RAM accesses **204**, **206**

determine that the read common misses both the main storage element **214** and the victim storage element **214**, the cache controller **220** sends an operation to the main storage element **214** to create a victim (e.g., to determine a line in the main storage element **214** to evict). In some examples, problems may arise when the main storage element **214** evicts a line that corresponds to outstanding writes. For example, a line in the main storage element **214** that is currently having data written to it by the store queue **212** can have outstanding writes located in any of the latches **402a-d**. In some examples, an error may occur when a new address replaces the line in the main storage element **214** that was evicted and the store queue **212** is still attempting to write data from a previous write operation into the same line.

(298) Examples disclosed herein include a read-invalidate operation to be issued to the store queue **212** when the cache controller **220** determines that a read operation is a read miss. The cache controller **220** sends the read-invalidate operation to the store queue **212** when the main storage element **214** creates a victim. The read-invalidate operation notifies the store queue **212** to compare the address of the victim to all the addresses stored in the latches **402a-d**. If any of the addresses in the latches **402a-d** match the address of the victim, the priority multiplexer **418** sends the data in the matching address to the latch **402e** to transmit the data to MUX circuit **314**. Such an operation may be referred to as data forwarding of in-flight data. For example, when the priority multiplexer **418** sends data corresponding to the victim to the MUX circuit **314** via the latch **402e**, the MUX circuit **314** forwards the data of the victim to the victim storage element **218** and/or the L2 cache **112**. Additionally, the read-invalidate operation invalidates the data stored in the latches **402a-d** when the latches **402a-d** include the address corresponding to the victim. For example, each of the latches **402a-d** include a valid tag. The valid tag indicates to the latches **402a-d** whether data is to be stored in the main storage element **214**. When the read-invalidate operation invalidates the valid tag, the latches **402a-d** discontinue writing the invalid data to main storage element **214**.

(299) In an example operation, the CPU **102** issues a write operation. In some examples, the write operation corresponds to an Address A. For example: Write operation=Write (Address A): Data 1 // write Data 1 to cache line corresponding to Address A

(300) The store queue **212** obtains the write operation and begins the process of writing data to Address A of the main store element **214**. In some examples, the Address A corresponds to cache line **1** in the main storage element **214**. Subsequently, the CPU **102** issues a read operation. In some examples, the read operation corresponds to an Address B. For example: Read operation=Read (Address B) // read Data 2 in cache line corresponding to Address B

(301) The example tag RAM accesses **204**, **206** check the addresses in the tag RAM **208** and the tag RAM **210** and determine neither of the tag RAMs **208**, **210** include Address B. The cache controller **220** obtains the read-miss result from the tag RAM accesses **204**, **206** and issues an allocate operation to determine a victim in the main storage element **214** to evict to the victim storage element **218** and/or the L2 cache **112**. When the cache controller **220** obtains a response corresponding to a victim, the cache controller **220** issues a read-invalidate operation to the store queue **212**.

(302) The read port **424** of the store queue **212** obtains the read-invalidate operation and obtains the address of the victim. The read port **424** sends the address of the victim to the data store **416** to be compared to all of the addresses stored in the latches **402a-d**. If the data store **416** determines any of the addresses stored in the latches **402a-d** match the address of the victim, the data store **416** outputs an operation to the priority multiplexer **418** to send the data corresponding to the victim address to the latch **402e**. The latch **402e** forwards the data to the MUX circuit **314** to send to the victim storage element **218** and/or the L2 cache **112**.

(303) Additionally, in response to the addresses stored in the latches **402a-d** matching the victim address, the latches **402a-d** invalidate the data corresponding to the victim address. For example, the valid tag in the latches **402a-d** are updated to indicate that the data in the latch corresponding to the victim address is invalid. In this example, the store queue **212** will not attempt to write that data

to the cache line corresponding to the previous address that was evicted.

(304) Example methods, apparatus, systems, and articles of manufacture for inflight data forwarding and invalidation are disclosed herein. Further examples and combinations thereof include the following:

(305) Example 1 includes an apparatus comprising a cache storage, a cache controller coupled to the cache storage and operable to receive a first memory operation, determine that the first memory operation corresponds to a read miss in the cache storage, determine a victim address in the cache storage to evict in response to the read miss, issue a read-invalidate command that specifies the victim address, compare the victim address to a set of addresses associated with a set of memory operations being processed by the cache controller, and in response to the victim address matching a first address of the set of addresses corresponding to a second memory operation of the set of memory operations, provide data associated with the second memory operation.

(306) Example 2 includes the apparatus of example 1, wherein the cache controller is to issue the read-invalidate command when a tag random access memory access issues a read-miss.

(307) Example 3 includes the apparatus of example 1, wherein the cache controller is to determine the victim address to evict from the cache storage to make room for a data of a third memory operation.

(308) Example 4 includes the apparatus of example 1, wherein the cache controller is operable to discontinue writing data associated with the second memory operation in the cache storage.

(309) Example 5 includes the apparatus of example 1, further including a first latch, a second latch, a third latch, and a fourth latch, wherein the first latch, the second latch, the third latch, or the fourth latch is to store data for the victim address.

(310) Example 6 includes the apparatus of example 5, wherein the first latch, the second latch, the third latch, and the fourth latch include a valid tag that is to be updated to invalid when any of the first latch, the second latch, the third latch, or the fourth latch store the data for the victim address.

(311) Example 7 includes the apparatus of example 1, further including a read port to initiate the comparison between the victim address and the set of addresses associated with the set of memory operations being processed by the cache controller when the read port obtains the read-invalidate command and the victim address.

(312) Example 8 includes a method to invalidate a write operation, the method comprising receiving a first memory operation, determining that the first memory operation corresponds to a read miss in a cache storage, determining a victim address in the cache storage to evict in response to the read miss, issuing a read-invalidate command that specifies the victim address, comparing the victim address to a set of addresses associated with a set of memory operations being processed by a cache controller, and in response to the victim address matching a first address of the set of addresses corresponding to a second memory operation of the set of memory operations, providing data associated with the second memory operation.

(313) Example 9 includes the method of example 8, further including issuing the read-invalidate command when a tag random access memory access issues a read-miss.

(314) Example 10 includes the method of example 8, further including determining the victim address to evict from the cache storage to make room for a data of a third memory operation.

(315) Example 11 includes the method of example 8, further including discontinuing writing the data associates with the second memory operation in the cache storage.

(316) Example 12 includes the method of example 8, further including storing data for the victim address in at least one of a first latch, a second latch, a third latch, and a fourth latch.

(317) Example 13 includes the method of example 12, further including updating a valid tag of the first latch, the second latch, the third latch, or the first latch to invalid when any of the first latch, the second latch, the third latch, and the fourth latch store the data for the victim address.

(318) Example 14 includes the method of example 8, further including initiating the comparison between the victim address and the set of addresses associated with the set of memory operations

being processed by the cache controller when a read port obtains the read-invalidate command and the victim address.

(319) Example 15 includes a system comprising a central processing unit to output a first memory operation, a cache coupled to the central processing unit, the cache including a cache storage, a cache controller coupled to the cache storage and operable to receive the first memory operation, determine that the first memory operation corresponds to a read miss in the cache storage, determine a victim address in the cache storage to evict in response to the read miss, issue a read-invalidate command that specifies the victim address, compare the victim address to a set of addresses associated with a set of memory operations being processed by the cache controller, and in response to the victim address matching a first address of the set of addresses corresponding to a second memory operation of the set of memory operations, provide data associated with the second memory operation and invalidate the second memory operation.

(320) Example 16 includes the system of example 15, wherein the cache is a first cache and wherein the first cache is operable to retrieve data associated with the first memory operation from a second cache in response to the cache storage not including the data associated with the first memory operation.

(321) Example 17 includes the system of example 15, wherein the cache includes a store queue to store write data of the second memory operation into the cache storage in multiple clock cycles, each clock cycle writing at least one portion of the data associated with the second memory operation to the cache storage.

(322) Example 18 includes the system of example 17, wherein the cache is a first cache and wherein the cache controller is operable to provide portions of data remaining in the store queue to a second cache for storing when the second memory operation is invalidated.

(323) Example 19 includes the system of example 15, wherein the cache includes a store queue to store outstanding writes, the second memory operation corresponding to an outstanding write.

(324) Example 20 includes the system of example 15, wherein the central processing unit outputs the first memory operation subsequently to outputting the second memory operation. FIG. 5 illustrates an example victim cache-based section of the level one (L1) cache **110** of FIGS. 1 and 3. FIG. 5 includes the tag RAM **210**, the victim storage **218**, the MESI RAM **300**, the example hit/miss comparison logic **306a**, **306b**, **306c**, the example replacement policy **308** (e.g., the replacement policy component **308**), the example address encoders **326a**, **326b**, **326c**, example multiplexers (MUXs) **330a**, **330b**, **330c**, the example address reads **332a**, **332b**, **332c**, an example scalar interface **502**, an example vector interface **504**, an example snoop address component **506**, and an example response multiplexer **508**. FIG. 5 illustrates a victim cache system that supports dual data paths and a coherency pipeline.

(325) In FIG. 5, the example scalar interface **502** is an interface coupling the L1 cache **110** of the data cache **108** of FIG. 1 to the example processing core **104a**. In some examples, the scalar interface **502** is an interface corresponding to a first data path (DP0) in the dual data path victim cache system. In some examples, the scalar interface **502** is an interface corresponding to a second data path (DP1) in the dual data path cache system. The example scalar interface **502** is a 64-bit wide bidirectional and/or unidirectional interface. Alternatively, the example scalar interface **502** may support a different quantity of bits (e.g., 32 bits, 128 bits, etc.). In some examples, the scalar interface **502** receives and/or obtains data from the core **104a**. In some examples, the scalar interface **502** sends data from the victim storage **218** to the core **104a**. The example scalar interface **502** is coupled to the example tag RAM **210**, and comparison logic **306** to compare an address from the CPU **102** to addresses in the tag RAM **210**.

(326) In FIG. 5, the example vector interface **504** is an interface coupling the L1 cache **110** of the data cache **108** of FIG. 1 to the example processing core **104a**. In some examples, the vector interface **504** is an interface corresponding to a second data path (DP1) in the dual data path cache system. In some examples, the vector interface **504** is an interface corresponding to a first data path

(DP0) in the dual data path cache system. The example vector interface **504** is a 512-bit wide interface. Alternatively, the example vector interface **504** may support a different quantity of bits (e.g., 256 bits, 1024 bits, etc.). In some examples, the vector interface **504** receives and/or obtains data from the core **104a**. In some examples, the vector interface **504** sends data from the victim storage **218** to the core **104a**. The example vector interface **504** is coupled to the example tag RAM **210**, the snoop address component **506**, and comparison logic **306b** to compare an address from the CPU **102** to addresses from the tag RAM **210**. In some examples, the scalar interface **502** and the vector interface **504** are implemented by the CPU interface **202** (FIG. 2). For example, the scalar interface **502** and the vector interface **504** can be included in the CPU interface **202**.

(327) In FIG. 5, the example tag RAM **210** is connected to the example interfaces **502**, **504**. However, in some examples, the tag RAM **210** may be coupled to the interfaces **502**, **504** via the example victim tag RAM access **206** (FIG. 2). The example tag RAM **210** is a dual-port fully associated tag RAM to support the dual data paths (e.g., scalar data path and vector data path). As used herein, a fully associated tag RAM **210** permits data from the main storage **214** (FIG. 2) to be stored in any cache block rather than forcing each memory address into one particular block. In some examples, because the tag RAM **210** is fully associated, the victim storage **218** is also fully associated.

(328) In FIG. 5, the comparison logic **306a** is connected to the scalar interface **502**, the tag RAM **210**, the address encoder **326a**, and the replacement policy component **308**. The comparison logic **306b** is connected to the vector interface **504**, the tag RAM **210**, the address encoder **326b**, and the replacement policy component **308**. The comparison logic **306a**, **306b** supports both scalar and vector data paths. For example, the comparison logic **306** can compare scalar data to data of tag RAM **210** and/or vector data to the data of the tag RAM **210**. The comparison logic **306** provides a hit or miss output to the replacement policy component **308** and the address encoders **326a**, **326b**. In some examples, the comparison logic **306** provides the hit or miss output to the multiplexers **330a**, **330b**.

(329) In FIG. 5, the example replacement policy component **308** is coupled to the comparison logic **306a**, **306b**. The example replacement policy component **308** is control/decision making logic. The example replacement policy component **308** dictates the entries (e.g., the data) of the example victim storage **218** based on a plurality of inputs. For example, the replacement policy component **308** can determine whether the cache controller **220** (FIG. 2) is to remove and/or enter entries to/from the victim storage **218**. The control logic of the replacement policy component **308** is configured to resolve address conflicts between the 2 addresses (e.g., scalar and vector) in such a way that data-consistency is maintained. FIG. 6 illustrates the control logic of the example replacement policy component **308**.

(330) In FIG. 5, the example address encoders **326a**, **326b** encode memory addresses into a form interpretable by the victim storage **218**. For example, the address encoders **326a**, **326b** can operate similarly to the address encoder **326** of FIGS. 3A-3D. The example address encoders **326a**, **326b** output encoded addresses to the multiplexers **330a**, **330b**. For example, address encoder **326a** can output encoded memory addresses to multiplexer **330a** and address encoder **326b** can output encoded memory addresses to multiplexer **330b**.

(331) In FIG. 5, the multiplexers **330a**, **330b** select an address in response to executing the comparison logic **306a**, **306b**. For example, the multiplexer **330a** can obtain an input from the comparison logic **306a** corresponding to a hit and therefore select the encoded address from the address encoder **326a**. Additionally and/or alternatively, the multiplexer **330a** may obtain an input from the comparison logic **306a** corresponding to a miss and therefore select the memory address chosen by the replacement policy component **308**. The example multiplexers **330a**, **330b** are coupled to address reads **332a**, **332b**.

(332) In FIG. 5, the example address reads **332a**, **332b** read the address values selected by the multiplexers into the victim storage **218**. For example, the address reads **332a**, **332b** can identify

addresses that are to be stored and/or removed from the victim storage **218**.

(333) In FIG. 5, the example victim storage **218** is a fully associative cache. For example, the fully associated victim storage **218** can place data, when data is fetched (e.g., victimized) from the main storage **214**, in any unused block of the cache. The placement of the data in the victim storage **218** is based on the replacement policy component **308**. For example, the replacement policy component **308** can determine when and where a line of data from the main storage **214** should be placed in the victim storage **218**. In some examples, when the address reads **332a**, **332b** read data into the victim storage **218**, the victim storage **218** outputs a response. In some examples, the victim storage **218** generates a DP0 response, a DP1 response, and a snoop response to be transmitted to the L2 data cache **112**. In some examples, the responses include data requested by the CPU **102**. For example, based on the address reads **332a**, **332b**, the victim storage **218** can output data from a memory address location if the CPU **102** has requested such data. Additionally and/or alternatively, based on the address reads **332a**, **332b**, the victim storage **218** can output eviction data to the L2 cache **112** if the replacement policy component **308** is replacing data in the victim storage **218**.

(334) In FIG. 5, the example snoop address component **506** is implemented by a snoop data path and/or otherwise interface. In addition to the two CPU data paths (e.g., DP0 and DP1), the L1 cache **110** includes the snoop data path to add coherency to the L1 cache **110**. The example snoop address component **506** is connected to the tag RAM **210** and comparison logic **306c**. In some examples, the snoop address component **506** obtains an example snoop request address issued by a higher-level data cache (e.g., the L2 data cache **112**) that issues an address read to the tag RAM **210**. For example, the snoop address component **506** attempts to read a memory address from the tag RAM **210**. Additionally, the snoop address component **506** provides the address read to the comparison logic **306c**. The snoop address component **506** and more generally, the coherency pipeline, is utilized to store the MESI state of every cache line in the victim storage **218** in the MESI RAM **300**. By storing the MESI state of every cache line in the MESI RAM **300**, the victim cache system supports coherency.

(335) In FIG. 5, the example MESI RAM **300** is connected to the comparison logic **306c**. The example MESI RAM **300** tracks the states of a cache line in the main storage **214** and the victim storage **218** by marking the cache line as modified, exclusive, shared, or invalid based on the comparison logic **306c**. For example, when the snoop address component **506** issues a read to the tag RAM **210**, the tag RAM **210** returns or does not return the correct memory address. Thus, the comparison logic **306c** compares the memory address of the snoop issued read and the tag RAM **210** results to determine the state of the cache line.

(336) In the example topology illustrated in FIG. 5, the response multiplexer **508** is coupled to the victim storage **218** and to the victim cache store queue **216**. In this manner, the response multiplexer **508** is configured to select between the data obtained by the victim storage **214** and/or the data obtained by the victim cache store queue **216** to transmit the data most recent to the system. Example operational description of the priority multiplexer **508** is described below.

(337) Methods and Apparatus for Multi-Banked Victim Cache with Dual Datapath

(338) In some victim caches, there is not support for a multi-bank architecture. For instance, a victim cache may be a monolithic storage device that does not support multiple accesses by a processor (e.g., a CPU) during the same clock cycle. For example, a request to access data in a single victim cache can lock up the entire victim cache. In such examples, there is a single register file capable of supporting one full victim cache line access per clock cycle. In some such examples, an entire victim cache line associated with the single victim storage can be locked to service the request because the single register file is allocated to the victim storage data bank that received such a request.

(339) Examples disclosed herein include a multi-banked victim cache with a dual datapath. In some disclosed examples, the victim storage **218** of FIG. 2 has a plurality of banks and, thus, is multi-

banked, or has a multi-bank data or memory structure. Advantageously, the multi-bank structure of the victim storage **218** can effectuate support for two or more accesses (e.g., CPU accesses) per clock cycle. Advantageously, the multi-bank structure of the victim storage **218** allows the two or more accesses to be serviced and/or fulfilled without performance overhead.

(340) In some disclosed examples, the victim storage **218** has a two-port fully associative tag ram (e.g., the tag ram **210** of FIG. 2). In some disclosed examples, hit/miss comparison is supported for both scalar and vector addresses. In such disclosed examples, based on the hit/miss decision on each of the datapaths (e.g., the scalar data path, the vector data path, etc.) the victim storage **218** is locked. In an example where an access to the victim storage **218** is a hit, then the location of the hit entry is selected. In an example, where an access to the victim storage **218** is a miss, then allocation policy is checked for a particular set of the victim storage **218**.

(341) In some disclosed examples, in a subsequent pipe stage, the selected sets of the victim storage **218** for both addresses are read and response(s) is/are given back to the CPU. In some disclosed examples, the set selection logic resolves address conflict between the two addresses (e.g., the scalar address and the vector address) in such a manner that data-consistency is maintained. In some disclosed examples, the control logic resolves address conflict as well.

(342) In some disclosed examples, bank arbitration logic calculates and/or otherwise determines the banks required for each of the addresses and selects the read on each of the banks. In some disclosed examples, the control logic resolves any bank conflict between the two addresses to effectuate serialization of the two memory operations. In some disclosed examples where both datapaths access the same address, the control logic can effectuate not issuing multiple allocation requests for the same address. Once the data is read from the victim storage **218**, example data muxing logic can fold the data into the correct data lanes before returning the data back to the CPU as read response(s).

(343) FIG. 7A is a schematic illustration of a first example encapsulated data cache system **700**. The first encapsulated data cache system **700** can be an example circuit implementation of the L1 cache **110** of FIG. 1 or portion(s) thereof, and/or, more generally, the data cache **108** of FIG. 1 or portion(s) thereof. The first encapsulated data cache system **700** is encapsulated to provide a unified storage view to an external system (e.g., one or more CPUs, one or more processors, external hardware, etc.). For example, the first encapsulated data cache system **700** can be encapsulated using firmware, software, hardware logic, etc., and/or a combination thereof to encapsulate and/or otherwise abstract all complexity related to read after write (RAW) hazards, write after write (WAW) hazards, data forwarding, etc. Advantageously, the unified architecture depicted in the example of FIG. 7A can provide separate ports to an external system (e.g., the CPU **102**) to simplify obtaining access to the most recent versions of data of interest.

(344) In the example of FIG. 7A, the first encapsulated data cache system **700** includes the victim cache store queue **216** of FIG. 2 and the victim storage **218** of FIG. 2. In FIG. 7A, the victim cache store queue **216** includes a plurality of first example independently addressable banks (e.g., data banks, store banks, data store banks, etc.) **702**. For example, the first independently addressable banks **702** can be bank queues, bank sub-queues, etc. In such examples, the bank queues, the bank sub-queues, etc., can include a plurality of storage elements, as depicted below in FIG. 7C and/or FIG. 7E. The first independently addressable banks **702** are victim cache store queue banks (e.g., victim cache store queue data banks, victim cache bank queues, victim cache bank sub-queues, etc.). In FIG. 7A, the victim cache store queue **216** includes 16 of the victim cache store queue banks **702**. Alternatively, there may be fewer or more than 16 instances of the victim cache store queue banks **702**.

(345) In FIG. 7A, the victim storage **218** includes a plurality of second example independently addressable banks (e.g., data banks, store banks, data store banks, etc.) **704**. The second independently addressable banks **704** are victim storage banks (e.g., victim storage data banks). In FIG. 7A, the victim storage **218** includes 16 of the victim storage banks **704**. Alternatively, there

may be fewer or more than 16 instances of the victim storage banks **704**. For example, the victim cache store queue **216** can include a plurality of bank queues or bank sub-queues including a first bank queue or a first bank sub-queue (e.g., VICTIM CACHE STORE QUEUE: BANK 1) and the victim storage **218** can include a plurality of data banks including a first data bank (e.g., VICTIM STORAGE: BANK 1).

(346) In some examples, the victim cache store queue **216** and the victim storage **218** each have 256 rows with each row having a line width of 1024 bits, but being divided into 16 banks as illustrated in FIG. 7A, with 64 bits per row in a given bank. In some examples, the victim cache store queue **216** and the victim storage **218** each have 384 rows with each row having a line width of 1024 bits, but being divided into 16 banks as illustrated in FIG. 7A, with 64 bits per row in a given bank. Of course, the cache size described in the above examples are only two examples and the disclosure is not limited to any particular cache line width, number of banks, or rows, etc. As depicted in FIG. 7A, each of the first banks **702** correspond to one of the second banks **704**. For example, a first bank (BANK 1) of the victim cache store queue **216** is in circuit with a first bank (BANK 1) of the victim storage **218** and, thus, corresponds with the first bank of the victim storage **218**.

(347) Advantageously, the multi-bank structure of the victim cache store queue **216**, the victim storage **218**, and/or, more generally, the first encapsulated data cache system **700**, can service read and write operations that are sent to the banks in parallel. In some examples, each bank arbitrates its own processes in response to the read and/or write operations. By allowing each bank of the first encapsulated data cache system **700** to operate independently, operation of the first encapsulated data cache system **700** is more efficient since an entire cache line is not locked up when a request is received. Rather, only the portion of the cache line allocated to the bank that received such a request would be locked.

(348) FIG. 7C depicts an example victim cache multi-bank structure **720**. In some examples, the L1 cache **110**, the L2 cache **112**, and/or the L3 cache **114** of FIG. 1 can have the victim cache bank structure **720**. In some examples, the victim cache store queue **216** of FIG. 2 and/or the victim storage **218** of FIG. 2 can have the victim cache multi-bank structure **720**. For example, the victim cache multi-bank structure **720** can be an example implementation of the victim cache store queue **216** and/or the victim storage **218**.

(349) The victim cache multi-bank structure **720** is a data or memory structure that includes 16 example banks (Banks 0-15) **722**, with each of the banks **722** having a data width of 64 bytes (e.g., bytes 0-7). Each of the banks **722** includes 16 example rows (Rows 0-15) **724**, with each of the rows **724** having a data width of 128 bytes (e.g., 128 bytes=8 bytes/bank*16 banks). In FIG. 7C, each of the banks **722** can store 128 bytes (e.g., 128 bytes=8 bytes/row*16 rows). In FIG. 7C, the victim cache multi-bank structure **720** can store 2048 bytes (e.g., 2048 bytes=128 bytes/bank*16 banks). The victim cache multi-bank structure **720** is independently addressable by bank. For example, the first row of the rows **724** has a starting row address of 0 and an ending row address of 127, a second row of the rows **724** has a starting row address of 128 and an ending row address of 255, etc.

(350) In some examples, a cache line can be 128 bytes of data that fits in a width of memory (e.g., DRAM) or storage unit (e.g., the main storage **214**, the victim storage **218**, etc.). In the example of FIG. 7C, a cache line can consume an entire row of the victim cache bank structure **720**. For example, a cache line can use one of the rows **724** of 16 banks, where each bank is 8 bytes wide. Advantageously, the victim cache bank structure **720** can enable 16 different cache lines to access data stored therein.

(351) FIG. 8A illustrates a schematic illustration of example victim cache tag (VCT) random access memory (RAM) **800**. The VCT RAM **800** can be an example implementation of the tag ram **210** of FIG. 2. For example, the VCT RAM **800** can store addresses of data stored in the victim cache store queue **216**, the victim storage **218**, etc., of FIG. 2. The VCT RAM **800** is a multi-bank VCT

RAM. For example, the VCT RAM **800** can include a plurality of banks (e.g., data banks, memory banks, etc.), such as 16 banks, although the VCT RAM **800** can have a different quantity of banks.

(352) The VCT RAM **800** includes example allocation ports **802**, **804**, **806** including a first example allocation port (AP0) **802**, a second example allocation port (AP1) **804**, and a third example allocation port (AP2) **806**. The VCT RAM **800** includes example read ports **808**, **810**, **812** including a first example read port (RP0) **808**, a second example read port (RP1) **810**, and a third example read port (RP2) **812**. The VCT RAM **800** includes an example LRU read port **814**. The VCT RAM **800** includes example output ports **816**, **818**, **820**, **822** including a first example output port (OP0) **816**, a second example output port (OP1) **818**, a third example output port (OP2) **820**, and a fourth example output port (OP3) **822**. Alternatively, VCT RAM **800** may fewer or more allocation ports, read ports, LRU read ports, and/or output ports than depicted in FIG. **8A**.

(353) The VCT RAM **800** includes the allocation ports **802**, **804**, **806** to fill the victim storage **218**. For example, the allocation ports **802**, **804**, **806** can be configured to receive requests from external hardware (e.g., the CPU **102**, the main storage **214**, etc.) to store data in the victim storage **218**. The first allocation port **802** is configured to receive first example data **824**. For example, the first allocation port **802** can receive the first data **824** from the write state machine associated with the scalar data path (DP0). The first data **824** includes WRM_TAG_UPDATE0 data, which can be representative of a control signal generated from the CPU interface **202** of FIG. **2** (e.g., the scalar data path (DP0)). The control signal can be generated to inform the VCT RAM **800** that the CPU interface **202** has a cache line to be moved from the CPU interface **202** to the victim storage **218** and, thus, the CPU interface **202** has an address to be moved from the CPU interface **202** to the tag ram **210**. The first data **824** includes VTAG_WR_TAG0 data, which can be representative of an address (e.g., a tag address) of the VCT RAM **800** that can correspond to an address of data to be stored in the victim cache **218**. The first data **824** includes VTAG_WR_SET0 data, which can be representative of the address of the victim cache **218** of where to store the data (e.g., the victim cache tag for DP0).

(354) The second allocation port **804** is configured to receive second example data **826**. For example, the second allocation port **804** can receive the second data **826** from the write state machine associated with the vector data path (DP1). The second data **826** includes WRM_TAG_UPDATE1 data, which can be representative of a control signal generated from the CPU interface **202** of FIG. **2** (e.g., the vector data path (DP1)). The control signal can be generated to inform the VCT RAM **800** that the CPU interface **202** has a cache line to be moved from the CPU interface **202** to the victim storage **218** and, thus, the CPU interface **202** has an address to be moved from the CPU interface **202** to the tag ram **210**. The second data **826** includes VTAG_WR_TAG1 data, which can be representative of an address (e.g., a tag address) of the VCT RAM **800** that can correspond to an address of data to be stored in the victim cache **218**. The second data **826** includes VTAG_WR_SET1 data, which can be representative of the address of the victim cache **218** of where to store the data (e.g., the victim cache tag for DP1).

(355) The third allocation port **806** is configured to receive third example data **828**. The third data **828** includes ARB_EVT_TAG_UPDATE data, which can be representative of a control signal generated from the main storage **214**. The control signal is an arbitration (ARB) evict (EVT) tag update control signal, which can be generated to inform the VCT RAM **800** that the main storage **214** has a cache line to be moved from the main storage **214** to the victim storage **218** and, thus, the main storage **214** has an address to be moved from the tag ram **208** to the tag ram **210**. The third data **828** includes ADP_EVT_WR_TAG data, which can be representative of an address (e.g., a tag address) of the VCT RAM **800** that can correspond to an address of data to be stored in the victim cache **218**. The third data **828** includes ADP_EVT_WR_SET data, which can be representative of the address of the victim cache **218** of where to store the data (e.g., the victim cache tag for the line moved from the main cache to the victim cache). For example, ADP_EVT_WR_TAG and ADP_EVT_WR_SET data can be referred to as address datapath (ADP) data. The first data **824**,

the second data **826**, and/or the third data **828** can be one or more data packets, one or more signals based on a communication protocol (e.g., an inter-integrated circuit (I2C) protocol), etc.

(356) The VCT RAM **800** includes the read ports **808**, **810**, **812** to provide interfaces for which external hardware (e.g., the CPU **102**) can request to read and/or otherwise access data stored in the victim storage **218**. The first read port **808** is configured to receive fourth example data **830**. For example, the first read port **808** can receive the fourth data **830** from the scalar interface **502** of the CPU **102**. The fourth data **830** includes ADP_ADDR_E2_DP0 data, which can be representative of an address of the victim storage **218** that the scalar interface **502** requests access to.

(357) The second read port **810** is configured to receive fifth example data **832**. For example, the second read port **810** can receive the fifth data **832** from the vector interface **504** of the CPU **102**. The fifth data **832** includes ADP_ADDR_E2_DP1 data, which can be representative of an address of the victim storage **218** that the vector interface **504** requests access to.

(358) The third read port **812** is configured to receive sixth example data **834**. For example, the third read port **812** can receive the sixth data **834** from the snoop address component **506** of FIG. 5 (e.g., the snoop interface). The sixth data **834** includes SNP_ADDR_E2_DP0 data, which can be representative of an address of the victim storage **218** that the snoop address component **506** requests access to.

(359) The LRU read port **814** is configured to receive seventh example data **836**. For example, the LRU read port **814** can receive the seventh data **836** from the replacement policy component **308** of FIGS. 3A-3D. The seventh data **836** includes LRU_SET_DP0 and LRU_SET_DP1, which can be representative of respective addresses associated with the least randomly used (LRU) cache lines of the victim storage **218**. For example, the LRU read port **814** can be a victim least randomly used (VLRU) read port configured to receive LRU data from the replacement policy component **308**.

(360) The VCT RAM **800** includes the output ports **816**, **818**, **820**, **822** to transmit outputs to external hardware (e.g., the CPU **102**, the main storage **214**, etc.) in response to a read request or a write request (e.g., an allocation request) associated with the victim storage **218**. The first output port **816** is configured to transmit first example output data **838**. For example, the first output port **816** can transmit the first output data **838** to the scalar interface **502**. The first output data **838** includes VTAG_HIT_DP0 data, which can indicate that data requested by the scalar interface **502** is stored in the victim storage **218**. The first output data **838** includes VTAG_MISS_DP0 data, which can indicate that the data requested by the scalar interface **502** is not stored in the victim storage **218**. The first output data **838** includes VTAG_SET_DP0 data, which can be representative of the address in the victim storage **218** where the data requested by the scalar interface **502** is stored.

(361) The second output port **818** is configured to transmit second example output data **840**. For example, the second output port **818** can transmit the second output data **840** to the vector interface **504**. The second output data **840** includes VTAG_HIT_DP1 data, which can indicate that data requested by the vector interface **504** is stored in the victim storage **218**. The second output data **840** includes VTAG_MISS_DP1 data, which can indicate that the data requested by the vector interface **504** is not stored in the victim storage **218**. The second output data **840** includes VTAG_SET_DP1 data, which can be representative of the address in the victim storage **218** where the data requested by the vector interface **504** is stored.

(362) The third output port **820** is configured to transmit third example output data **842**. For example, the third output port **820** can transmit the third output data **842** to the replacement policy component **308**. The third output data **842** includes VTAG_ADDR_DP0 data, which can be representative of the address of the victim storage **218** that is associated with the request LRU_SET_DP0 by the replacement policy component **308** received at the LRU read port **814**.

(363) The fourth output port **822** is configured to transmit fourth example output data **844**. For example, the fourth output port **822** can transmit the fourth output data **844** to the replacement policy component **308**. The fourth output data **844** includes VTAG_ADDR_DP1 data, which can

be representative of the address of the victim storage **218** that is associated with the request LRU_SET_DP1 by the replacement policy component **308** received at the LRU read port **814**.

(364) FIG. **8B** illustrates another schematic illustration of the example implementation of the VCT RAM **800** of FIG. **8A**. For example, FIG. **8B** can correspond to the victim cache-based section of the L1 cache **110** of FIG. **5**, or portion(s) thereof.

(365) FIG. **8B** depicts another example implementation of the VCT RAM **800** of FIG. **8A**. In FIG. **8B**, the VCT RAM **800** includes 16 example sets **846**. For example, the sets **846** include a first set VCT_ADDR[0], a second set VCT_ADDR[1], etc. Alternatively, the VCT RAM **800** may include fewer or more than 16 of the sets **846**.

(366) The VCT RAM **800** is victim cache tag storage configured to store addresses (e.g., tag addresses) that correspond to the sets **846**. Each of the sets **846** is coupled to a respective one of first example comparators **850** and a respective one of second example comparators **852**. The first comparators **850** can be an example implementation of the comparison logic **306a** of FIGS. **3** and/or **5**. The second comparators **852** can be an example implementation of the comparison logic **306b** of FIGS. **3** and/or **5**. The first set VCT_ADDR[0] is coupled to a first one of the first comparators **850** and a first one of the second comparators **852**, the second set VCT_ADDR[1] is coupled to a second one of the first comparators **850** and a second one of the second comparators **852**, etc.

(367) The first comparators **850** and the second comparators **852** are coupled to respective example address encoder logic circuits **854**, **856** including a first example address encoder logic circuit **854** and a second example address encoder logic circuit **856**. The first comparators **850** are coupled to the first address encoder logic circuit **854**. The second comparators **852** are coupled to the second address encoder logic circuit **856**.

(368) The address encoder logic circuit **854**, **856** can be example implementations of the address encoders **326a**, **326b** of FIGS. **3** and/or **5**. For example, the first address encoder logic circuit **854** can be an example implementation of the address encoder **326a**. In other examples, the second address encoder logic circuit **856** can be an example implementation of the address encoder **326b**.

(369) A first example multiplexer **858A** is coupled to the first address encoder logic circuit **854**. The first address encoder logic circuit **854** includes a first example decoder **860A**, a first example inverter (e.g., a logic gate, an inverting logic gate, etc.) **862A**, a first example AND gate (e.g., a logic gate, an AND logic gate, etc.) **864A**, a first example OR gate (e.g., a logic gate, an OR logic gate, etc.) **866A**, a first example encoder **868A**, a third example comparator **870A**, and a fourth example comparator **872A**.

(370) FIG. **8B** depicts only one instance each of the first AND gate **864A**, the first OR gate **866A**, the third comparator **870A**, and the fourth comparator **872A** for clarity. However, each of the first AND gate **864A**, the first OR gate **866A**, the third comparator **870A**, and the fourth comparator **872A** can be representative of 16 instances of said component. For example, the first AND gate **864A** can be representative and/or otherwise implement 16 instances of the first AND gate **864A**, the first OR gate **866A** can be representative of and/or otherwise implement 16 instances of the first OR gate **866A**, etc.

(371) The first multiplexer **858A** has a first input to receive ADP_ADDR_E2_DP0, which is representative of an address requested by the DP0 interface from the E2 Arbitration stage of FIGS. **3A-3D**. The first multiplexer **858A** has a second input to receive SNP_ADDR_E2_DP0, which is representative of a snoop address requested by the snoop interface from the E2 Arbitration stage of FIGS. **3A-3D**. The first multiplexer **858A** has a select input to receive SNP_ADDR_EN_DP0, which is representative of an enable signal from the snoop interface that, when asserted, can invoke the first multiplexer **858A** to select the second input. An output of the first multiplexer **858A** is coupled to inputs of the first comparators **850**. For example, each of the first comparators **850** can compare an address from the output of the first multiplexer **858A** to a corresponding one of the sets of the VCT RAM **800**. In such examples, a first one of the first comparators **850** can compare an

address included in ADP_ADDR_E2_DP0 to the address stored in VCT_ADDR[0]. In some such examples, the first one of the first comparators **850** can output and/or otherwise generate a logic one in response to the address from the first multiplexer **858A** matching the address stored in VCT_ADDR[0].

(372) Outputs of the first multiplexer **858A** are coupled to a first input of the third comparator **870A**. An output of an example DP0 read finite-state machine (FSM) (READ_FSM_DP0) **873** and/or an output of an example DP0 write finite-state machine (WRITE_FSM_DP0) **874** is coupled to a second input of the third comparator **870A**. The DP0 read finite-state machine **873** and the DP0 write finite-state machine **874** are hardware implemented finite-state machines that execute logic on data from the scalar interface **502** of FIG. 5. For example, the DP0 read finite-state machine **873** and/or the DP0 write finite-state machine **874** can transmit one or more of WRM_TAG_UPDATE0, VTAG_WR_TAG0, and VTAG_WR_SET0 to the third comparator **870A** and/or the fourth comparator **872A**. The DP0 read finite-state machine **873** and/or the DP0 write finite-state machine **874** can transmit VTAG_WR_TAG0 to a sixth example comparator **872B** included in the second address encoder logic circuit **856**. WRM_TAG_UPDATE0, VTAG_WR_TAG0, and VTAG_WR_SET0 are described above in connection with FIG. 8A.

(373) The first decoder **860A** is a 4×16 decoder. The first decoder **860A** has an input to receive VTAG_WR_SET0 data, which can be representative of an in-flight address from the scalar interface **502** to the victim storage **218**. The first decoder **860A** has an output coupled to an input of the first inverter **862A**. The first decoder **860A** can convert the in-flight address to a bit vector where each bit is inverted by one of the 16 instances of the first inverter **862A**. An output of the first inverter **862A** is coupled to a first input of the first AND gate **864A**. A second input of the first AND gate **864A** is coupled to the result bit of the tag comparison from the first comparators **850** with set 0 (e.g., VCT_ADDR[0]) and the output of the first multiplexer **858A**. For example, the second input of the first AND gate **864A** can be configured to receive HIT_DP0 data, which can be representative of a 16-bit vector, where each of the bits can correspond to whether the ADP_ADDR_E2_DP0 data is a hit (e.g., a bit value of 1) or a miss (e.g., a bit value of 0) in the victim storage **218**.

(374) An output of the first AND gate **864A** is coupled to a first input of the first OR gate **866A**. An output of the third comparator **870A** is coupled to a second input of the first OR gate **866A**. An output of the fourth comparator **872A** is coupled to a third input of the first OR gate **866A**. An output of the first OR gate **866A** is coupled to an input of the first encoder **868A**. The first encoder **868A** is a 16×4 encoder. For example, the first encoder **868A** can generate HIT_ADDR0 data, which can be representative of VTAG_SET_DP0 of FIG. 8A. In such examples, HIT_ADDR0 can correspond to the first output data **838** of FIG. 8A.

(375) The second address encoder logic circuit **856** includes a second example decoder **860B**, a second example inverter (e.g., a logic gate, an inverting logic gate, etc.) **862B**, a second example AND gate (e.g., a logic gate, an AND logic gate, etc.) **864B**, a second example OR gate (e.g., a logic gate, an OR logic gate, etc.) **866B**, a second example encoder **868B**, a fifth example comparator **870B**, and the sixth example comparator **872B**.

(376) FIG. 8B depicts only one instance each of the second AND gate **864B**, the second OR gate **866B**, the fifth comparator **870B**, and the sixth comparator **872B** for clarity. However, each of the second AND gate **864B**, the second OR gate **866B**, the fifth comparator **870B**, and the sixth comparator **872B** can be representative of 16 instances of said component. For example, the second AND gate **864B** can be representative and/or otherwise implement 16 instances of the second AND gate **864B**, the second OR gate **866B** can be representative of and/or otherwise implement 16 instances of the second OR gate **866B**, etc.

(377) The second address encoder logic circuit **856** has a first input to receive ADP_ADDR_E2_DP1, which is representative of an address requested by the DP1 interface from the E2 Arbitration stage of FIGS. 3A-3D. ADP_ADDR_E2_DP1 is coupled to inputs of the second

comparators **852**. For example, each of the second comparators **852** can compare the address included in ADP_ADDR_E2_DP1 to a corresponding one of the sets of the VCT RAM **800**. In such examples, a first one of the second comparators **852** can compare an address included in ADP_ADDR_E2_DP1 to the address stored in VCT_ADDR[0]. In some such examples, the first one of the second comparators **852** can output and/or otherwise generate a logic one in response to the address included in ADP_ADDR_E2_DP1 matching the address stored in VCT_ADDR[0]. (378) ADP_ADDR_E2_DP1 is coupled to a first input of the fifth comparator **870B**. An output of an example DP1 read finite-state machine (READ_FSM_DP1) **875** and/or an output of an example DP1 write finite-state machine (WRITE_FSM_DP1) **876** is coupled to a second input of the fifth comparator **870B**. The DP1 read finite-state machine **875** and the DP1 write finite-state machine **876** are hardware implemented finite-state machines that execute logic on data from the vector interface **504** of FIG. 5. For example, the DP1 read finite-state machine **875** and/or the DP1 write finite-state machine **876** can transmit one or more of WRM_TAG_UPDATE1, VTAG_WR_TAG1, and VTAG_WR_SET1 to the fifth comparator **870B** and/or the sixth comparator **872B**. The DP1 read finite-state machine **875** and/or the DP1 write finite-state machine **876** can transmit VTAG_WR_TAG1 to the fourth comparator **872A** included in the first address encoder logic circuit **854**. WRM_TAG_UPDATE1, VTAG_WR_TAG1, and VTAG_WR_SET1 are described above in connection with FIG. 8A.

(379) The second decoder **860B** is a 4×16 decoder. The second decoder **860B** has an input to receive VTAG_WR_SET1 data, which can be representative of an in-flight address from the vector interface **504** to the victim storage **218**. The second decoder **860B** has an output coupled to an input of the second inverter **862B**. The second decoder **860B** can convert the in-flight address to a bit vector where each bit is inverted by one of the 16 instances of the second inverter **862B**. An output of the second inverter **862B** is coupled to a first input of the second AND gate **864B**. A second input of the second AND gate **864B** is coupled to the result bit of the tag comparison from the second comparators **852** with set 0 (e.g., VCT_ADDR[0]) and ADP_ADDR_E2_DP1. For example, the second input of the second AND gate **864B** can be configured to receive HIT_DP1 data, which can be representative of a 16-bit vector, where each of the bits can correspond to whether the ADP_ADDR_E2_DP1 data is a hit (e.g., a bit value of 1) or a miss (e.g., a bit value of 0) in the victim storage **218**.

(380) An output of the second AND gate **864B** is coupled to a first input of the second OR gate **866B**. An output of the fifth comparator **870B** is coupled to a second input of the second OR gate **866B**. An output of the sixth comparator **872B** is coupled to a third input of the second OR gate **866B**. An output of the second OR gate **866B** is coupled to an input of the second encoder **868B**. The second encoder **868B** is a 16×4 encoder. For example, the second encoder **868B** can generate HIT_ADDR1 data, which can be representative of VTAG_SET_DP1 of FIG. 8A. In such examples, HIT_ADDR1 can correspond to the second output data **840** of FIG. 8A.

(381) In example operating conditions, the first address encoder logic circuit **854** can receive a first read address (e.g., ADP_ADDR_E2_DP0) for a store instruction requested by the CPU **102** of FIG. 1, where the first read address is received from the E2 pipeline stage (ARBITRATION) of FIGS. 3A-3D. In response to the first multiplexer **858A** selecting the first input, the first multiplexer **858A** can transmit the first read address to the first comparators **850**. The first comparators **850** can compare the read address to the respective ones of the sets **846**. If, a first one of the first comparators **850** determines that the read address matches an address stored at the corresponding one of the sets **846**, the first one of the first comparators **850** can assert and/or otherwise output a logic one, otherwise the first one of the first comparators **850** outputs a logic zero. Accordingly, the first comparators **850** can generate HIT_DP0 on a 16-bit data bus as a 16-bit vector in response to the comparisons. For example, the 16-bit vector can be an implementation of a train (e.g., a pulse train) of signals representative of a bit vector. The first comparators **850** can transmit HIT_DP0 to the first AND gate **864A**.

(382) Portion(s) of the first address encoder logic circuit **854** and/or the second address encoder logic circuit **856** are cache hit-miss conversion logic. For example, the first AND gate **864A**, the third comparator **870A**, and/or the fourth comparator **872A** can be an example implementation of cache hit-miss conversion logic. In such examples, the second AND gate **864B**, the fifth comparator **870B**, and/or the sixth comparator **872B** can be an example implementation of the cache hit-miss conversion logic.

(383) The first AND gate **864A** is configured to determine whether to convert a hit (e.g., a cache hit) to a miss (e.g., a cache miss). For example, a new address (VTAG_WR_SET0) can be written to the victim storage **218** during the E3 pipeline stage of FIGS. **3A-3D** while the first read address is being processed during the E2 pipeline stage. In such examples, the first decoder **860A** can convert the new address into a 16-bit value, which can be inverted by the first inverter **862A**. The first inverter **862A** can transmit the inverted 16-bit value to the first input of the first AND gate **864A**. The first AND gate **864A** can assert a logic one in response to VTAG_WR_SET0 not matching the address in HIT_DP0 and, thus, does not convert the cache hit to a cache miss. In other examples, the first AND gate **864A** can output a logic zero in response to VTAG_WR_SET0 matching the address in HIT_DP0 and, thus, converts the cache hit to a cache miss because the address requested in ADP_ADDR_E2_DP0 has been overwritten and is no longer available at that address.

(384) In example operating conditions, the third comparator **870A** and the fifth comparator **870B** can be configured to convert a cache miss to a cache hit. For example, the third comparator **870A** can determine that the first read address (ADP_ADDR_E2_DP0) in the VCT RAM **800** requested during the E2 pipeline stage is getting written in the E3 pipeline stage, which is represented by VTAG_WR_TAG0. In such examples, the third comparator **870A** can assert a logic one in response to ADP_ADDR_E2_DP0 matching VTAG_WR_TAG0 and, thus, convert the cache miss to a cache hit and HIT_ADDR0 can be updated with VTAG_WR_SET0 because the data will be available when the ADP_ADDR_E2_DP0 address is read during the E3 pipeline stage.

(385) In example operating conditions, the fourth comparator **872A** and the sixth comparator **872B** can be configured to convert a cache miss to a cache hit. For example, the fourth comparator **872A** can determine that the first read address (ADP_ADDR_E2_DP0) in the VCT RAM **800** requested during the E2 pipeline stage is getting written in the E3 pipeline stage by the vector interface **504**, which is represented by VTAG_WR_TAG1. In such examples, the fourth comparator **872A** can assert a logic one in response to ADP_ADDR_E2_DP0 matching VTAG_WR_TAG1 and, thus, convert the cache miss to a cache hit and HIT_ADDR0 can be updated with VTAG_WR_SET1 because the data will be available when the ADP_ADDR_E2_DP0 address is read during the E3 pipeline stage.

(386) In example operating conditions, the first OR gate **866A** and the second OR gate **866B** can be configured to generate an output to a corresponding one of the first encoder **868A** or the second encoder **868B**. For example, the first OR gate **866B** can transmit a 16-bit vector representative of a cache miss (e.g., 16 bit values of 0) or a cache hit (e.g., 16-bit value of an address of the cache hit). In such examples, the first encoder **868A** can encode the 16-bit value from the first OR gate **866A** as a 4-bit address and, thus, generate HIT_ADDR0. Such example operations can be applicable to the second OR gate **866B**, the second encoder **868B**, and/or, more generally, the second address encoder logic circuit **856**.

(387) Example methods, apparatus, systems, and articles of manufacture for multi-banked victim cache with dual datapath are disclosed herein. Further examples and combinations thereof include the following:

(388) Example 1 includes a cache system comprising a storage element that includes banks operable to store data, ports operable to receive memory operations in parallel, wherein each of the memory operations has a respective address, and a plurality of comparators coupled such that each of the comparators is coupled to a respective port of the ports and a respective bank of the banks

and is operable to determine whether a respective address of a respective memory operation received by the respective port corresponds to the data stored in the respective bank.

(389) Example 2 includes the cache system of example 1, further including victim cache including victim cache data, and the data is tag data corresponding to the victim cache data.

(390) Example 3 includes the cache system of example 1, wherein the memory operations include a first memory operation, further including cache including cache data, the data is tag data corresponding to the cache data, and a first comparator of the plurality of the comparators to generate a signal in response to detecting a cache hit based on first data associated with the first memory operation corresponding to the data in the storage element.

(391) Example 4 includes the cache system of example 1, wherein the plurality of the comparators is a plurality of first comparators, and further including an address encoder logic circuit including a decoder having a decoder input and a decoder output, the decoder input coupled to a first finite-state machine associated with a scalar interface of a central processing unit, a plurality of inverters having a respective inverter input and a respective inverter output, the inverter inputs coupled to the decoder output, a plurality of AND logic gates having a respective first input and a respective second input, a respective one of the first inputs coupled to a respective one of the inverter outputs, a respective one of the second inputs coupled to a respective one of outputs of the plurality of the first comparators, a plurality of second comparators coupled to respective ones of the outputs of the plurality of the first comparators and the first finite-state machine, and a plurality of third comparators coupled to the respective outputs of the plurality of the first comparators and a second finite-state machine associated with a vector interface of the central processing unit.

(392) Example 5 includes the cache system of example 4, wherein a first AND logic gate of the plurality of the AND logic gates is to output a logic low signal to convert a cache hit to a cache miss in response to a first address associated with the scalar interface matching a second address associated with the cache hit.

(393) Example 6 includes the cache system of example 4, wherein a first comparator of the plurality of the second comparators is to output a logic high signal to convert a cache miss to a cache hit in response to a first address associated with the scalar interface matching a second address associated with the cache miss.

(394) Example 7 includes the cache system of example 4, wherein a first comparator of the plurality of the third comparators is to output a logic high signal to convert a cache miss to a cache hit in response to a first address associated with the vector interface matching a second address associated with the cache miss.

(395) Example 8 includes the cache system of example 4, wherein the address encoder logic circuit includes a plurality of OR logic gates having a respective third input, a respective fourth input, and a respective fifth input, the third inputs coupled to a respective one of outputs of the plurality of the AND logic gates, the fourth inputs coupled to a respective one of outputs of the plurality of the second comparators, the fifth inputs coupled to a respective one of outputs of the plurality of the third comparators, and an encoder having an encoder input and an encoder output, the encoder input coupled to outputs of the plurality of the OR logic gates, the encoder output coupled to an input of a multiplexer.

(396) Example 9 includes the cache system of example 1, wherein the plurality of the comparators is a plurality of first comparators, and further including an address encoder logic circuit including a decoder having a decoder input and a decoder output, the decoder input coupled to a first finite-state machine associated with a vector interface of a central processing unit, a plurality of inverters having a respective inverter input and a respective inverter output, the inverter inputs coupled to the decoder output, a plurality of AND logic gates having a respective first input and a respective second input, a respective one of the first inputs coupled to a respective one of the inverter outputs, a respective one of the second inputs coupled to a respective one of outputs of the plurality of the first comparators, a plurality of second comparators coupled to respective ones of the outputs of the

plurality of the first comparators and the first finite-state machine, and a plurality of third comparators coupled to the respective outputs of the plurality of the first comparators and a second finite-state machine associated with a vector interface of the central processing unit.

(397) Example 10 includes a victim cache system comprising a victim cache store queue including first data banks, and a victim storage coupled to the victim cache store queue, the victim storage including second data banks different from the first data banks.

(398) Example 11 includes the victim cache system of example 10, wherein at least one of the first data banks or the second data banks includes 16 data banks.

(399) Example 12 includes the victim cache system of example 10, wherein the first data banks includes a first data bank having 8 addresses each having a data width of 1 byte, the first data bank configured to store 64 bits.

(400) Example 13 includes the victim cache system of example 10, wherein the second data banks includes a second data bank having 8 addresses each having a data width of 1 byte, the second data bank configured to store 64 bits.

(401) Example 14 includes the victim cache system of example 10, wherein a respective one of the first data banks is coupled to a respective one of the second data banks.

(402) Example 15 includes a method to identify a cache hit in a victim cache, the method comprising receiving a first read address from a first interface of a processor, receiving a second read address from a second interface of the processor, comparing the first read address and the second read address to data banks of a multi-bank victim cache tag memory, in response to mapping at least one of the first read address or the second read address to a first data bank of the data banks, identifying a cache hit, and outputting a cache hit address representative of a victim cache address of cache data stored in the victim cache.

(403) Example 16 includes the method of example 15, wherein the first interface is a scalar interface and the second interface is a vector interface.

(404) Example 17 includes the method of example 15, further including in response to identifying the cache hit, comparing the cache hit address to a first address received from a finite-state machine associated with the first interface, the first address representative of an address written to in the victim cache in a previous pipeline stage, and in response to the cache hit address matching the first address, converting the cache hit to a cache miss.

(405) Example 18 includes the method of example 15, further including in response to identifying the cache hit, comparing the cache hit address to a first address received from a finite-state machine associated with the second interface, the first address representative of an address written to in the victim cache in a previous pipeline stage, and in response to the cache hit address matching the first address, converting the cache hit to a cache miss.

(406) Example 19 includes the method of example 15, further including in response to not mapping at least one of the first read address or the second read address to any of the data banks, identifying a cache miss, comparing the first read address to a first address received from a finite-state machine associated with the first interface, the first address representative of an address written to in the victim cache in a subsequent pipeline stage, and in response to the first read address matching the first address, converting the cache miss to the cache hit.

(407) Example 20 includes the method of example 15, further including in response to not mapping at least one of the first read address or the second read address to any of the data banks, identifying a cache miss, comparing the second read address to a first address received from a finite-state machine associated with the second interface, the first address representative of an address written to in the victim cache in a subsequent pipeline stage, and in response to the second read address matching the first address, converting the cache miss to the cache hit.

(408) Example 21 includes the method of example 15, wherein comparing the first read address and the second read address to the data banks of the multi-bank victim cache tag memory is executed substantially in parallel.

(409) Methods and Apparatus for Allocation in a Victim Cache System

(410) Data cache architectures including a victim cache system enable the main cache (e.g., the main storage **214**) to allocate data to a victim cache (e.g., the victim storage **218**) when the main cache needs to create a victim. For example, when an incoming instruction is a read-miss, the main cache needs to make room for data of the read-miss. The main cache may be a direct mapped cache such that the read-miss can only be stored in one location, indicated by the address of the read-miss. The main cache may allocate data of the read-miss location to be moved to the victim cache when the data is dirty and evict data of the read-miss location to be sent out to higher level memory locations when the data of the location is clean. For example, the main cache controller may elect a modified (e.g., dirty) line in the main storage to be victimized and an unmodified (e.g., clean) line in the main storage to be evicted. A dirty and/or modified line of data is data that has been updated by the CPU **102**. For example, when the cache line in the main storage **214** is present only in the current state and has been modified from the value in the main memory (e.g., so that a read of the data in extended memory **106** is flagged or permitted because it is not up-to-date or valid), the cache line is dirty and/or modified and therefore is a candidate for allocation (e.g., victimization). The example MESI RAM **300** stores and/or tracks the cache line states (e.g., clean, dirty, shared, exclusive, etc.), therefore allowing and/or enabling the identification of dirty lines in the main storage **214**.

(411) An allocation policy of the main storage may instruct the main cache controller to elect to victimize a modified line because the data for the memory address is not located in higher level cache or is located in higher level cache but is outdated. Such an allocation policy may instruct the main cache controller to not allocate/victimize a clean and/or shared line in the main storage because that line includes data at the memory address that is already located in the higher level cache (e.g., L2 cache, L3 cache, extended memory, etc.).

(412) However, such an allocation policy creates latency (e.g., increased the time it would take for the CPU to retrieve the requested data) when only allocating dirty and/or modified lines in the L1 cache **110**. The latency is a result of using extra clock cycles to retrieve from higher level memory. For example, due to the parallel connection of the main storage **214** and the victim storage **218**, retrieving data from the higher level memories takes more time than retrieving data from the victim storage **218**.

(413) In examples disclosed herein, the allocation policy reduces the latency of allocating cache lines from the main cache (e.g., the main storage **214**) to the victim cache (e.g., the victim storage **218**). For example, due to the architecture of the L1 cache **110** (e.g., the parallel connection between the main storage **214** and the victim storage **218**), the allocation policy enables the main cache controller **222** to allocate any cache line (e.g., clean and/or dirty) in the main storage **214** to the victim storage **218**. The configuring of the main storage **214** and the victim storage **218** in parallel enables the victim storage **218** to act as an extension of the main storage **214**. Thus, in addition to providing a location to store victim lines, the victim storage **218** can also provide extra associativity to the main storage **214** by moving any cache line (e.g., clean and/or dirty) from the main storage **214** to the victim storage **214**. For example, because the victim storage **218** and main storage **214** are accessed in parallel (e.g., at the same time) by the CPU **102**, if a clean line was to be accessed and it was allocated to the victim storage **218**, the latency for accessing the clean line of data would not be effected.

(414) In an example operation, the main cache controller **222** obtains an instruction (e.g., a memory operation) from the CPU interface **202** (e.g., the scalar interface **502** and/or the vector interface **504**). The main cache controller **222** obtains results from the main tag RAM access **204** corresponding to the instruction. For example, the main tag RAM access **204** determines if the address in the instruction matches an address in the tag RAM **208** and provides the results to the main cache controller **222**. In some examples, the main cache controller **222** determines that the main storage **214** is to victimize and/or allocate data to the victim storage **218** based on the

instruction and the results from the main tag RAM access **204**. For example, the main cache controller **222** determines if the instruction is a read operation and if the read operation is a miss (e.g., determined based on the main tag RAM access **204** results). If the read operation is a miss, the main cache controller **222** determines that the main storage **214** needs to allocate the line, way, block, slot, etc. of data for allocation in the victim storage **218**.

(415) In some examples, the main cache controller **222** allocates the line corresponding to the address of the read-miss operation. For example, if a read operation wants to read data from Address A but the main storage **214** stores different data for that address location (e.g., Address A'), the main cache controller **222** has to remove the data of Address A' because the data of Address A can only be stored in one location (e.g., Address A) of the main storage **214**. The main cache controller **222** identifies the location in the direct mapped cache (e.g., the main storage **214**) that is mapped to Address A' and allocates data of Address A' to the victim storage **218**. Therefore, the read operation can store data of Address A in the main storage **214**.

(416) Example methods, apparatus, systems, and articles of manufacture for allocation of data are disclosed herein. Further examples and combinations thereof include the following:

(417) Example 1 includes an apparatus comprising a first cache storage, a second cache storage, a cache controller coupled to the first cache storage and the second cache storage and operable to receive a memory operation that specifies an address, determine, based on the address, that the memory operation evicts a first set of data from the first cache storage, determine that the first set of data is unmodified relative to an extended memory, and cause the first set of data to be stored in the second cache storage.

(418) Example 2 includes the apparatus of example 1, wherein the cache controller is to identify a portion in the first cache storage corresponding to a directly mapped location of the address.

(419) Example 3 includes the apparatus of example 1, wherein the cache controller is to allocate the first set of data to the second cache storage to create space for data of a second memory operation.

(420) Example 4 includes the apparatus of example 1, wherein the cache controller is to receive a second memory operation that specifies a second address, determine, based on the second address, that the second memory operation evicts a second set of data from the first cache storage, determine that the second set of data is modified relative to an extended memory, and cause the second set of data to be stored in the second cache storage.

(421) Example 5 includes the apparatus of example 1, wherein the cache controller is to perform an outstanding transaction of a store queue on the first set of data during the eviction.

(422) Example 6 includes the apparatus of example 1, wherein the second cache storage is to provide associativity to the first cache storage when the second cache storage stores the first set of data.

(423) Example 7 includes the apparatus of example 1, wherein the second cache storage is a victim cache storage.

(424) Example 8 includes a method comprising receiving a memory operation that specifies an address, determining, based on the address, that the memory operation evicts a first set of data from a first cache storage, determining that the first set of data is unmodified relative to an extended memory, and cause the first set of data to be stored in a second cache storage.

(425) Example 9 includes the method of example 8, further including identifying a portion in the first cache storage corresponding to a directly mapped location of the address.

(426) Example 10 includes the method of example 8, further including allocating the first set of data to the second cache storage to create space for data of a second memory operation.

(427) Example 11 includes the method of example 8, further including receiving a second memory operation that specifies a second address, determining, based on the second address, that the second memory operation evicts a second set of data from the first cache storage, determining that the second set of data is modified relative to an extended memory, and causing the second set of data to

be stored in the second cache storage.

(428) Example 12 includes the method of example 8, further including performing an outstanding transaction of a store queue on the first set of data during the eviction.

(429) Example 13 includes the method of example 8, wherein storing the first set of data allocated from the first cache storage includes providing associativity to the first cache storage.

(430) Example 14 includes the method of example 8, wherein the second cache storage is a victim cache storage.

(431) Example 15 includes a system comprising a central processing unit to output a memory operation that specifies an address, and a cache coupled to the central processing unit, the cache operable to obtain the memory operation from the central processing unit, the cache further including a first cache storage, a second cache storage, and a cache controller to determine, based on the address of the memory operation, that the memory operation evicts a first set of data from the first cache storage, determine that the first set of data is unmodified relative to an extended memory, and cause the first set of data to be stored in the second cache storage.

(432) Example 16 includes the system of example 15, wherein the cache further includes a first tag random access memory coupled to the central processing unit, the first cache storage, and the cache controller, and a second tag random access memory coupled to the central processing unit, the second cache storage, and the cache controller.

(433) Example 17 includes the system of example 15, wherein the cache controller is to determine, based on a second address of a second memory operation, that the second memory operation evicts a second set of data from the first cache storage, determine that the second set of data is modified relative to an extended memory, and cause the second set of data to be stored in the second cache storage.

(434) Example 18 includes the system of example 15, wherein the cache controller is to perform an outstanding transaction of a store queue on the first set of data during the eviction.

(435) Example 19 includes the system of example 15, wherein the first cache storage and the second cache storage are connected in parallel to the central processing unit.

(436) Example 20 includes the system of example 15, wherein the first cache storage is a direct mapped cache and the second cache storage is a fully associative cache, the direct mapped cache coupled and the fully associative cache coupled to the central processing unit in parallel.

(437) Methods and Apparatus to Facilitate Read-Modify-Write Support in a Coherent Victim Cache with Parallel Data Paths

(438) To facilitate a memory system that operates with memory coherence, that is, to facilitate a memory system in which corresponding memory locations for each processing element contain the same stored data, some example systems are prevented from including a victim storage. Examples disclosed herein overcome this challenge by facilitating a fully coherent memory system in which the victim storage (e.g., the victim storage **218**) can operate in memory coherence operations in the system.

(439) To provide coherency on the data cache **108**, the example scalar interface **502** (e.g., a sixty-four-bit wide interface) and the example vector interface **504** (e.g., a five hundred and twelve-bit wide interface) are included in the victim storage **218** of the L1 cache **110**. In this manner, the victim storage **218** of the L1 cache **110** is capable of servicing two data paths from the CPU **102**.

(440) In some examples disclosed herein, the snoop address component **506** is utilized to store the MESI state of every cache line in the victim storage **218** in the MESI RAM **300**. By storing the MESI state of every cache line in the MESI RAM **300**, the victim cache system supports coherency.

(441) Additionally, the example snoop address component **506** obtains an example snoop request from a higher-level memory cache memory (e.g., the L2 data cache **112**) and is transmits such a snoop request to the tag RAM **210** and comparison logic **306c**. In some examples, the snoop address component **506** triggers an address read instruction to the tag RAM **210**. For example,

when the L2 data cache **112** is unable to identify data requested in a read request from the CPU **102**, the L2 data cache **112** transmits the snoop address component **506** to the lower level, L1 data cache **110** to attempt to execute a read request of a memory address in the tag RAM **210**. Additionally, the snoop address component **506** provides the address read to the comparison logic **306c**. In the event a miss is identified (e.g., the read issued to the tag RAM **210** returns a miss), an appropriate snoop response indicating that a read was unsuccessful is generated and transmitted to the higher-level data cache that generated the snoop address component **506** (e.g., the L2 data cache **112**).

(442) Alternatively, in the event a hit is identified (e.g., the read issued to the tag RAM **210** returns a hit), the state of the corresponding line in the MESI RAM **300** is read and the example address encoder **326c** generates an address value for use by the victim storage **218** in obtaining the data. Accordingly, an appropriate response indicating that the read request was successful is generated and transmitted back to the higher-level data cache that generated the snoop address component **506** (e.g., the L2 data cache **112**) as the snoop response. In examples disclosed herein, the snoop response includes the data in the victim storage **218** corresponding to the address in the tag RAM **210**.

(443) In an example operation, the corresponding store queue (e.g., the victim cache store queue **216**) may be processing a write instruction to the address that is being read via the snoop address. Accordingly, while the victim storage **218** is servicing a snoop request (e.g., while the snoop request is being processed in response to the snoop address component **506** obtaining the snoop request), the victim cache store queue **216** forwards the data from the victim cache store queue **216** (e.g., the data stored in latch **402e**) to the response multiplexer **508**. In this manner, any state change obtained by the vector interface **504** due to the snoop address and any recently updated address obtained from the victim cache store queue **216** is forwarded to the higher-level data cache (e.g., the L2 data cache **112**).

(444) In the topology illustrated in FIG. 5, the coherency pipeline is longer than the victim cache pipeline to provide enough time for the victim cache controller **224** to properly order a potential snoop response and/or subsequent CPU **102** operation in the event such a snoop response and/or subsequent CPU **102** operation is issued to a higher level memory controller.

(445) In examples disclosed herein, the victim storage **218** of the L1 data cache **110** is capable of issuing tag-updates to higher level cache controller in the event tracking of cache lines is requested. In this manner, the victim storage **218** can facilitate tracking of cache lines to distinguish between exclusive and modified cache elements.

(446) In examples disclosed herein, the victim storage **218** of the L1 data cache **110** supports global coherence operation by allowing a global coherence operation to process on a fully associative cache and iterating over the tag RAM **208** based on sets.

(447) While the above-mentioned operation is described in connection with the L1 data cache **110**, any of the above-mentioned operations and/or elements may be implemented on any of the L2 data cache **112**, the L3 data cache **114**, and/or any additional level data cache in the data cache **108**.

(448) Example methods, apparatus, systems, and articles of manufacture to facilitate read-modify-write support in a coherent victim cache with parallel data paths are disclosed herein. Further examples and combinations thereof include the following:

(449) Example 1 includes an apparatus comprising a random-access memory configured to be coupled to a central processing unit via a first interface and a second interface, the random-access memory configured to obtain a read request indicating a first address to read via a snoop interface, an address encoder coupled to the random-access memory, the address encoder to, when the random-access memory indicates a hit of the read request, generate a second address corresponding to a victim cache based on the first address, and a multiplexer coupled to the victim cache to transmit a response including data obtained from the second address of the victim cache.

(450) Example 2 includes the apparatus of example 1, wherein the read request is transmitted by a

level two data cache.

(451) Example 3 includes the apparatus of example 1, wherein the multiplexer is configured to be coupled to a victim cache storage queue, the multiplexer to obtain input data from the victim cache storage queue indicating a value to be written to the second address.

(452) Example 4 includes the apparatus of example 3, wherein the multiplexer is configured to transmit the response including the input data from the victim cache storage queue.

(453) Example 5 includes the apparatus of example 1, wherein the random-access memory is a tag random access memory.

(454) Example 6 includes the apparatus of example 1, wherein the victim cache is a multi-bank memory.

(455) Example 7 includes the apparatus of example 1, wherein the first interface is a scalar interface and the second interface is a vector interface.

(456) Example 8 includes the apparatus of example 7, wherein the scalar interface is a sixty-four-bit wide interface, the vector interface is a five hundred and twelve-bit wide interface.

(457) Example 9 includes a system comprising a controller, a first memory, a second memory configured to obtain a snoop request from the first memory, the second memory to execute a read request of a random-access memory in response to the snoop request, when the random-access memory indicates a hit of the read request, generate an address corresponding to a victim cache, the address to be used in reading the victim cache, and transmit a response to the second memory including data obtained from the address of the victim cache.

(458) Example 10 includes the system of example 9, wherein the first memory is a level two cache memory and the second memory is a level one cache memory.

(459) Example 11 includes the system of example 9, wherein the second memory includes a multiplexer configured to be coupled to a victim cache storage queue, the multiplexer to obtain input data from the victim cache storage queue indicating a value to be written to the address of the victim cache in the second memory.

(460) Example 12 includes the system of example 11, wherein the multiplexer is configured to transmit the response including the input data from the victim cache storage queue.

(461) Example 13 includes the system of example 9, wherein the second memory is coupled to the controller via a scalar interface and a vector interface, and the second memory is coupled to the first memory via a snoop interface.

(462) Example 14 includes the system of example 13, wherein the scalar interface is a sixty-four-bit wide interface, the vector interface is a five hundred and twelve-bit wide interface.

(463) Example 15 includes the system of example 13, wherein the controller is a central processing unit.

(464) Example 16 includes a method comprising obtaining a read request indicating a first address to read via a snoop interface, when the read request is a hit, generate a second address corresponding to a victim cache based on the first address, and generating a response including data obtained from the second address of the victim cache.

(465) Example 17 includes the method of example 16, further including transmitting the read request from a level two data cache.

(466) Example 18 includes the method of example 16, further including obtaining input data from a victim cache storage queue indicating a value to be written to the second address.

(467) Example 19 includes the method of example 18, further including transmitting the response including the input data from the victim cache storage queue.

(468) Example 20 includes the method of example 16, wherein the victim cache is a multi-bank memory.

(469) Methods and Apparatus for Eviction in Dual Datapath Victim Cache System

(470) In a victim cache system, the main cache (e.g., the main storage **214**) victimizes (e.g., allocates) cache lines to the victim cache (e.g., victim storage **218**) when the main cache needs to

store new data. When the main cache creates a victim, the replacement policy (e.g., replacement policy component **308**) determines where the victim can be stored in the victim cache (e.g., the victim storage **218**). In some examples, the victim cache is full and thus needs to evict data to the higher level cache memories (e.g., L2 cache **112**, L3 cache **114**, extended memory **106**). The victim cache (e.g., victim storage **218**) also evicts data to the higher level cache memories when a write-miss occurs. For example, the victim storage **218** includes a write-miss buffer that buffers write-miss data. The replacement policy may utilize fixed schemes to determine what data to evict from the victim cache. For example, eviction schemes such as First In First Out (FIFO) scheme, Random scheme, and Least Recently Used (LRU) scheme. However, such eviction schemes are not configured to efficiently manage the eviction of data from the victim cache when there are two or more data paths.

(471) For example, the FIFO scheme evicts a block of data that has been in the victim cache for the longest time (e.g., the data in slot 0 if it is the first time the eviction is occurring). In a dual data path victim cache system (e.g., the victim cache section illustrated in FIG. 5), the FIFO scheme does not work when two incoming CPU accesses (e.g., instructions) are both misses and/or are both allocating instructions. In such an example, if two sets from the victim cache are not speculatively locked and read, then additional stalls are created.

(472) In a different example, the main storage **214** stores an address A and the victim storage **218** stores an address B. The CPU **102** sends two requests: the CPU **102** requests to read data of address A' utilizing DP0 and the CPU **102** requests to read data from address B utilizing DP1. The data of address A' is mapped to the same location as the address A in the main storage **214**. Therefore, the first data path DP0 needs to allocate the data in address A to the victim storage **218** (e.g., allocate data of address A' into the main storage **214**). In a regular FIFO scheme, address A would be allocated to slot 0 of the victim storage **218**. However, the slot 0 is occupied by address B and the CPU **102** is requesting to access the address B in parallel to requesting to allocate data of the read instruction to address A'. Therefore, if the data path DP0 evicted address B from slot 0, a stall (e.g., a cycle or a number of cycles for which the CPU **102** is stalled waiting for the memory access) would be created in order to retrieve the evicted data later from higher level caches.

(473) Examples disclosed herein include an eviction scheme that efficiently evicts data from the victim storage **218** in a dual data path victim cache system. The example eviction scheme is implemented by the replacement policy component **308** of FIGS. 3 and 5. For example, the replacement policy component **308** includes logic that analyzes inputs and produces an output. For example, the replacement policy component **308** obtains inputs from the hit-miss comparison logic **306a** and **306b** and determines what location in the victim storage **218** is to be replaced by data from the main storage **214** and/or replaced by data from the write-miss instruction. In examples disclosed herein, the replacement policy component **308** reserves entries (e.g., ways, sets, cache lines, etc.) in the victim storage **218** that correspond to eviction locations. For example, the replacement policy component **308** speculatively locks (e.g., reserves) first and second victim cache lines (e.g., sets) that are specifically for eviction. The first and second victim cache lines may be locked for specific data paths (e.g., first victim cache line locked for DP0 and second victim cache line locked for DP1). The eviction logic implemented by the replacement policy component **308** is described in further detail below in connection with FIG. 6.

(474) Turning to FIG. 6, an example first table **602** and an example second table **604** are depicted. The example first and second tables **602**, **604** describe the logic of the example replacement policy component **308**. The example replacement policy component **308** utilizes an LRU value. As used herein, an LRU value corresponds to a way that is the least recently used portion in the victim storage **218**. In some examples, the LRU value corresponds to a location in the victim storage element **218** that has not been recently accessed by the first or second data path. In a fully associated cache (e.g., the victim storage **218**), an address can be mapped to any "way" in the cache. Thus, a way is a block in the cache that stores data and includes a tag. In some LRU

schemes, when the victim cache needs to evict data, the data in the location indicated by the LRU value (e.g., the data in the portion not recently accessed in the victim storage **218**) is evicted. For example, in a 3-way cache, an address A, an address B, and an address C occupy all three ways (e.g., blocks, slots, etc.). In this example, address A has been recently accessed by the CPU **102**, address B was accessed just prior to the address A access, and address C was accessed prior to address B. Therefore, address C is the least recently used address and the portion mapped to address C has a location equal to the LRU value.

(475) The example first table **602** of FIG. **6** includes a first data path scenario **606**, a second data path scenario **608**, and a third data path scenario **610**. The example data path scenarios **606**, **608**, **610** correspond to the validity of the transactions (e.g., accesses) on DP0 and DP1 of FIG. **5**.

Validity of a transaction corresponds to whether the CPU **102** sends valid instructions to the L1 cache **110**. A valid instruction on the first data path DP0 and an invalid instruction on the second data path DP1 corresponds to when the processing cores send an instruction on the first data path DP0 and leave the second data path DP1 idle. In some examples, during a particular clock cycle, the first data path DP0 includes invalid instructions. In some examples, during a particular clock cycle, the second data path DP1 includes invalid instructions. Additionally and/or alternatively, both the data paths DP0 and DP1 can include valid instructions and/or can include invalid instructions.

(476) In the example first table **602**, the first data path scenario **606** corresponds to when both data paths (DP0, DP1) of FIG. **5** include valid transactions (DP0_valid=1, DP1_valid=1) in the same clock cycle. In some examples, the replacement policy component **308** obtains results and/or instructions from both data paths DP0, DP1.

(477) In the example first table **602**, the second data path scenario **608** corresponds to the first data path (DP0) including a valid transaction (DP0_valid=1) and the second data path (DP1) including an invalid transaction (DP1_valid=0). In such an example, the replacement policy component **308** determines that the transaction of DP1 includes invalid instructions.

(478) In the example first table **602**, the third data path scenario **610** corresponds to the first data path (DP0) including an invalid transaction (DP0_valid=0) and the second data path (DP1) including a valid transaction (DP1_valid=1). In some examples, the replacement policy component **308** determines that that transaction of DP0 includes invalid instructions.

(479) The example first data path scenario **606** in the example first table **602** includes an example hit-hit action **612**, an example miss-hit action **614**, an example hit-miss action **616**, and an example miss-miss action **618**. The example hit-hit action **612** is indicative that the address on DP0 and the address on DP1 hit (e.g., matches) an address in either of the tag RAMs **208**, **210**. In other words, the example hit-hit action **612** is indicative that the address on DP0 and the address on DP1 hit an address in main storage **214**, the victim storage **218** and/or the write miss cache (described above in connection with FIGS. **3A-3D**). The example miss-hit action **614** is indicative that the address on DP0 does not match an address in main storage **214**, the victim storage **218**, and/or the write miss cache and that the address on DP1 hits an address in the main storage **214**, the victim storage **218** and/or the write miss cache. In some examples, the miss portion of the miss-hit action **614** is indicative that DP0 is to allocate a cache line from the main storage **214** to the victim storage **218** and thus, is evicting data from the victim storage **218** to make room for the cache line of the main storage **214**. The example hit-miss action **616** is indicative that the address on DP0 hits an address in the victim storage **218** and/or the write miss cache and that the address on DP1 does not match an address in the victim storage **218** and/or the write miss cache. In some examples, the miss portion of the hit-miss action **616** is indicative that DP1 is to allocate a cache line from the main storage **214** to the victim storage **218** and thus, is evicting data from the victim storage **218** to make room for the cache line of the main storage **214**. The example miss-miss action **618** is indicative that the addresses of DP0 and DP1 both do not have a match in the victim storage **218** and/or the write miss cache. In some examples, the miss-miss action **618** is indicative that both data paths

DP0 and DP1 are to allocate data from the main storage **214** to the victim storage **218**. Additionally and/or alternatively, a miss is indicative that a write-miss is to write data to the write-miss buffer in the victim storage **218**.

(480) The example second data path scenario **608** in the example first table **602** includes an example DP0 hit action **620** and an example DP0 miss action **622**. In the second data path scenario **608**, the first data path DP0 includes a valid transaction and the second data path DP1 includes an invalid transaction. In some examples, the replacement policy component **308** disregards the second data path DP1 because the second data path DP1 does not include instructions.

(481) The example third data path scenario **610** in the example first table **602** includes an example DP1 hit action **624** and an example DP1 miss action **626**. The example DP1 hit action **624** corresponds to the action the replacement policy component **308** takes when the address of the second data path DP1 matches an address in the victim storage **218** and/or the write miss cache. The example DP1 miss action **624** corresponds to the action the replacement policy component **308** takes when the address of the second data path DP1 does not match an address in the victim storage **218** and/or the write miss cache. Similarly to the example DP0 hit action **620** and the DP0 miss action **622** under the example second data path scenario **608**, the example replacement policy component **308** disregards the transaction of the first data path DP0 in the example third data path scenario **610** because the transaction is invalid.

(482) In the example first table **602**, DP0 Hit Way indicates a portion of the victim storage **218** and/or the write miss cache that should be accessed (e.g., read from, evicted, written to, etc.) by the first data path DP0 when the instruction is a hit. In the example first table **602**, DP1 Hit Way is a portion in the victim storage **218** and/or the write miss cache that should be accessed by the second data path DP1 when the instructions is a hit.

(483) In the example first table **602**, the variable ‘Y’ is the variable that indicates the location of the current way selected as the LRU and indicates where the first data path DP0 should remove data from. In some examples, Y is assigned to the DP0 pointer. For example, when DP0 needs to evict a portion in the victim storage **218**, then DP0 pointer points to the location Y (e.g., the LRU way) for eviction. In some examples, the replacement policy component **308** is to store an indicator of the LRU way of the victim storage **218** to be replaced by DP0. For example, the replacement policy component **308** keeps an indicator, that can be accessed by the cache controller **220**, that a particular way, not recently accessed, is available for eviction by the first data path DP0. As used herein, the terms “pointer” and “indicator” may be used interchangeably.

(484) In the example first table **602**, the variable ‘Y+1’ indicates the location of the next LRU way where the second data path DP1 should remove data from. In some examples, ‘Y+1’ is assigned to the DP1 pointer. For example, when the second data path DP1 needs to evict a portion in the victim storage **218**, the DP1 pointer points to the location Y+1 (e.g., the next LRU way) for eviction. In some examples, the replacement policy component **308** is to store an indicator of the LRU way of the victim storage **218** to be replaced by DP1. For example, the replacement policy component **308** keeps an indicator, that can be accessed by the cache controller **220**, that a particular way, not recently accessed, is available for eviction by the first data path DP0.

(485) In a first example operation of the replacement policy component **308**, both data paths (DP0 and DP1) include valid transactions (e.g., indicated in first data path scenario **606**) and both data paths (DP0 and DP1) include addresses that match addresses in the example victim storage element **218** (e.g., the hit-hit action **612**). For example, the comparison logic **306a** and **306b** compares the addresses of DP0 and DP1 to addresses in the example tag RAM **210** and provides a hit result to the replacement policy component **308**. Since both data paths are valid and both data paths are hits, the DP0 Hit Way points to the portion (e.g., the way) in the victim storage **218** that includes hit/matching data corresponding to data path DP0 and DP1 Hit Way points to the way that includes the hit/matching data corresponding to data path DP1. The victim storage **218** responds with the requested data. The location of the LRU way does not increment and remains at location Y.

(486) In a second example operation of the replacement policy component **308**, both data paths (DP0 and DP1) include valid transactions (e.g., indicated in first data path scenario **606**), the first data path DP0 is a miss, and the second data path DP1 hits the victim storage **218** (e.g., the miss-hit action **614**). For example, the first comparison logic **306a** returns a “miss” result to the replacement policy component **308** and the second comparison logic **306b** returns a “hit” result to the replacement policy component **308**. The DP1 Hit Way points to the portion (e.g., way) in the victim storage **218** that includes the hit/matching data (e.g., the hit way). In some examples, the first data path DP0 is to evict data from a portion (e.g., way) in the victim storage **218**. Therefore, the DP0 pointer points to the location Y in the victim storage **218** that is to be evicted.

(487) Before eviction occurs, the replacement policy component **308** determines if the DP1 Hit Way matches the address of the LRU way (e.g., location Y). If the DP1 Hit Way does not have the same location as Y, the DP0 pointer points to the way having a location equal to the location Y (e.g., the LRU way). If the DP1 Hit Way does match the address of the LRU way, the DP0 pointer points to the location of the DP1 pointer (Y+1) so that the DP0 can evict data without conflicting with DP1 Hit Way.

(488) In the second example scenario **614**, the LRU value and the next LRU value are incremented based on which location was evicted. For example, if DP0 evicted data from location Y+1 (e.g., the DP1 Hit Way matches the location of the DP0 pointer), the LRU value is incremented twice and the next LRU value is incremented twice. Otherwise, if DP0 evicted data from location Y (e.g., DP1 Hit Way did not match the location of DP0 pointer) the LRU value is incremented once and the next LRU value is incremented once.

(489) In a third example operation of the replacement policy component **308**, both data paths (DP0 and DP1) include valid transactions (e.g., indicated in first data path scenario **606**), the first data path DP0 is a hit, and the second data path DP1 is a miss (e.g., the hit-miss action **616**). For example, the first comparison logic **306a** returns a “hit” result to the replacement policy component **308** and the second comparison logic **306b** returns a “miss” result to the replacement policy component **308**. The DP0 Hit Way points to the way in the victim storage **218** that includes the hit/matching data. In some examples, the miss causes the second data path DP1 to evict a way to make room in the victim storage **218**. Therefore, the DP1 pointer points to location Y+1 in the victim storage **218** that is to be evicted.

(490) Before eviction occurs, the replacement policy component **308** determines if the DP0 Hit Way matches the address of the next LRU way (e.g., location Y+1). If the replacement policy component **308** determines the DP0 Hit Way matches the address of the next LRU way (e.g., Y+1), the DP1 pointer points to the location of the DP0 pointer (e.g., location Y) so that the DP1 can evict data without conflicting with DP0 Hit Way. If the DP0 Hit Way does not match the address of the next LRU way, then the DP1 evicts data from location Y+1.

(491) In the third example scenario **616**, the LRU value and the next LRU value are incremented based on which location was evicted. For example, if DP1 evicted data from location Y (e.g., the DP0 Hit Way matches the location of the DP1 pointer), the LRU value is incremented once and the next LRU value is incremented once. Otherwise, if DP1 evicted data from location Y+1 (e.g., DP0 Hit Way did not match the location of DP1 pointer) the LRU value is incremented twice and the next LRU value is incremented twice.

(492) In a fourth example operation of the replacement policy component **308**, both data paths (DP0 and DP1) include valid transactions (e.g., indicated in first data path scenario **606**) and both data paths are flagged as misses (e.g., column **618**). For example, the comparison logic **306a** and **306b** returns “miss” results to the replacement policy component **308** when both addresses in the data paths DP0 and DP1 are not found and/or matched with the addresses in the tag RAMs **208**, **210**. In the fourth operation, both data paths DP0 and DP1 are to evict ways in the victim storage **218**. Therefore, DP0 pointer points to location Y and DP1 pointer points to location Y+1.

(493) When both eviction operations are complete, the LRU value is incremented by two (e.g.,

Y+2) and the next LRU value is incremented by two (e.g., (Y+1)+2). In a subsequent operation, if DP0 and DP1 are misses, DP0 Way points to the new LRU value (e.g., Y+2) and DP1 Way points to the next LRU value (e.g., (Y+1)+2).

(494) In a fifth example operation of the replacement policy component **308**, the first data path DP0 is a valid transaction and the second data path DP1 is an invalid transaction (e.g., indicated in second data path scenario **608**). In the fifth example operation, the first data path DP0 is a hit (e.g., indicated in the DP0 hit action **620**). For example, the comparison logic **306a** returns a “hit” result to the replacement policy component **308**. The DP0 Hit Way points to the way in the victim storage **218** that includes the matching data. The LRU value (Y) remains the same because no data is to be evicted in the clock cycle.

(495) In a sixth example operation of the replacement policy component **308**, the first data path DP0 is a valid transaction and the second data path DP1 is an invalid transaction (e.g., indicated in second data path scenario **608**). In the sixth example operation, the first data path DP0 is a miss (e.g., indicated in the DP0 miss action **622**). For example, the comparison logic **306a** returns a “miss” result to the replacement policy component **308**. In this example, the first data path DP0 is to evict data from the victim storage **218**. The example DP0 pointer points to the location Y (e.g., the LRU way). After eviction, the LRU value is incremented (e.g., Y+1).

(496) In a seventh example operation of the replacement policy component **308**, the first data path DP0 is an invalid transaction and the second data path DP1 is a valid transaction (e.g., indicated in third data path scenario **610**). In the seventh example operation, the second data path DP1 is a hit (e.g., indicated in the DP1 hit action **624**). For example, the comparison logic **306b** returns a “hit” result to the replacement policy component **308**. The DP1 Hit Way points to the way in the victim storage **218** that includes the matching data. The LRU value (Y) remains the same because no data is to be evicted in the clock cycle.

(497) In an eighth example operation of the replacement policy component **308**, the first data path DP0 is an invalid transaction and the second data path DP1 is a valid transaction (e.g., indicated in third data path scenario **610**). In the eighth example operation, the second data path DP1 is a miss (e.g., indicated in the DP1 miss action **626**). For example, the comparison logic **306b** returns a “miss” result to the replacement policy component **308**. In this example, the second data path DP1 is to evict data from the victim storage **218**. The DP1 pointer points to the location Y (e.g., the LRU way). The DP1 pointer does not point to location Y+1 because of the invalid transaction of DP0. Otherwise, when both transactions are valid, DP1 always points to Y+1 (e.g., unless switched when DP0 Hit Way matches Y+1). After the second data path DP1 evicts the data from location Y, the LRU value is incremented (e.g., Y+1).

(498) Turning to the second table **604**, second table **604** illustrates the incrementation of the LRU value when the first data path DP0 and/or the second data path DP1 is allocating data into the victim storage **218**. For example, when a read-miss occurs, the main storage **214** allocates a line of data to the victim storage **218** utilizing one of the data paths. The second table **604** includes a first valid column **626**, a second valid column **628**, a first allocate column **630**, a second allocate column **632**, a first LRU interference **634**, a second LRU interference **636**, and an LRU increment column **638**.

(499) In the second table **604**, the example first valid column **626** corresponds to the validity of the second data path transaction. For example, a zero (0) indicates the DP1 transaction is invalid and a one (1) indicates that the DP1 transaction is valid. The example second valid column **628** corresponds to the validity of the first data path transaction. For example, a zero (0) indicates the DP0 transaction is invalid and a one (1) indicates that the DP0 transaction is valid.

(500) In the second table **604**, the example first allocate column **630** indicates the allocation status of the second data path DP1. The allocation status corresponds to allocation of data from the main storage **214** to the victim storage **218** in a clock cycle. For example, a zero (0) indicates that the second data path DP1 is not allocating data into the victim storage **218** and a one (1) indicates that

the second data path DP1 is allocating data into the victim storage **218**. The example second allocate column **632** indicates the allocation status of the first data path DP0. For example, a zero (0) indicates that the first data path DP0 is not allocating data into the victim storage **218** and a one (1) indicates that the first data path DP0 is allocating data into the victim storage **218**. When a data path is allocating data into the victim storage **218**, the data path evicts a way (e.g., slot, block, etc.) to make room for the data being allocated. In some examples, data is allocated to the victim storage **218** when a read-miss occurs in the main storage **214**.

(501) In the second table **604**, the first LRU interference column **634** indicates whether the first data path DP0 hits the same location in the victim storage **218** as the location of the second data path allocate pointer. For example, the address of the first data path DP0 is located in the least recently used location of the victim storage **218**. The first LRU interference column **634** includes a one (1) to indicate that the first data path DP1 hit location equals the location of the second data path DP1 allocate pointer.

(502) The second LRU interference column **636** indicates whether the second data path DP1 hits the same location in the victim storage **218** as the location of the second data path allocate pointer. For example, the address of the second data path DP1 is located in the least recently used location of the victim storage **218**. The second LRU interference column **636** includes a one (1) to indicate that the second data path DP1 hit location equals the location of the first data path allocate pointer. As used herein, the first data path allocate pointer points to the location Y (LRU value) when DP0 is to allocate and the second data path allocate pointer points to the location Y+1 (next LRU value) when the DP1 is to allocate. In some examples, the pointers notify the cache controller **220** to evict a portion of the victim storage **218** to the higher level caches (e.g., L2 **112**, L3 **114**, extended memory **106**). The example replacement policy component **308** may initialize the first data path allocate pointer to point to location Y (LRU portion) and initialize the second data path allocate pointer to point to Y+1 (next LRU portion).

(503) In the second table **604**, the LRU increment column **628** indicates the incrementation of the LRU value, Y. For example, the replacement policy component **308** increments the LRU value by one (e.g., Y+1), by two (e.g., Y+2), or by nothing (e.g., Y). In some examples, the incrementation of the LRU value depends on the status of the data paths DP0 and DP1.

(504) In an example operation, both the first data path DP0 and the second data path DP1 include valid transactions. The example replacement policy component **308** determines if any of the data paths are allocating. For example, the cache controller **220** sends information to the replacement policy component **308** when the main storage **214** needs to allocate data.

(505) If the first data path DP0 is allocating data (e.g., moving data from the main storage **214** to the victim storage **218**), the first data path DP0 evicts data (e.g., indicated by the first data path allocate pointer) from the victim storage **214**. To evict data, the replacement policy component **308** determines whether the second data path DP1 was a hit and where the hit location is. For example, the replacement policy component **308** analyzes the location of the address of the second data path DP1 and determines if that location matches the location of the first data path allocate pointer.

(506) If the second data path hit location is equal to the location of the first data path allocate pointer (e.g., DP1 hit location=Y), the replacement policy component **308** updates the first data path allocate pointer to point to the next LRU value (Y+1) (e.g., notifies the cache controller **220** to evict data of next LRU value). In this manner, the second data path DP1 reads/writes from the hit location Y and the first data path DP0 evicts data of the LRU location Y+1. By way of this operation, the first data path DP0 does not evict the read/write data of DP1. After the eviction has occurred (e.g., eviction of data from Y+1 in the victim storage **218**), the replacement policy component **308** increments the first data path allocate pointer by two and the second data path allocate pointer by two. For example, the replacement policy component **308** increments LRU value (Y) by two and the next LRU value (Y+1) by two because DP0 just evicted location Y+1, and therefore, the new LRU value will be Y+2. This operation is illustrated at row **640**.

(507) If the second data path hit location is not equal to the location of the first data path allocated pointer (e.g., DP1 hit location does not equal Y), the replacement policy component **308** notifies the cache controller **220** that location Y is to be evicted. In this manner, the cache controller **220** evicts data from the location Y in the victim storage **218**. After eviction has occurred (e.g., eviction of data from Y in the victim storage **218**), the replacement policy component **308** increments the first data path allocate pointer by one and the second data path allocate pointer by one. For example, the replacement policy component **308** increments LRU value (Y) by one and the next LRU value (Y+1) by one because DP0 just evicted location Y, and therefore, the new LRU value will be Y+1. This operation is illustrated at row **642**.

(508) In a different operation, the second data path DP1 is allocating data (e.g., moving data from the main storage **214** to the victim storage) and the second data path DP1 evicts data (e.g., indicated by the second data path allocate pointer) from the victim storage **214**. To evict data, the replacement policy component **308** determines whether the first data path DP0 was a hit and where the hit location is. For example, the replacement policy component **308** analyzes the location of the address of the first data path DP0 and determines if that location matches the location of the second data path allocate pointer.

(509) If the first data path hit location is equal to the location of the second data path allocate pointer (e.g., DP0 hit location=Y+1), the replacement policy component **308** updates the second data path allocate pointer to point to the LRU value (Y) (e.g., notifies the cache controller **220** to evict data of LRU value). In this manner, the first data path DP0 reads/writes from the hit location Y+1 and the second data path DP1 evicts data of the LRU location Y. By way of this operation, the second data path DP1 does not evict the read/write data of DP0. After the eviction has occurred (e.g., eviction of data from Y in the victim storage **218**), the replacement policy component **308** increments the first data path allocate pointer by one and the second data path allocate pointer by one. For example, the replacement policy component **308** increments LRU value (Y) by one and the next LRU value (Y+1) by one because DP1 just evicted location Y, and therefore, the new LRU value will be Y+1. This operation is illustrated at row **644**.

(510) If the first data path hit location is not equal to the location of the second data path allocated pointer (e.g., DP0 hit location does not equal Y+1), the replacement policy component **308** notifies the cache controller **220** that location Y+1 is to be evicted. In this manner, the cache controller **220** evicts data from the location Y+1 in the victim storage **218**. After eviction has occurred (e.g., eviction of data from Y+1 in the victim storage **218**), the replacement policy component **308** increments the first data path allocate pointer by two and the second data path allocate pointer by two. For example, the replacement policy component **308** increments the LRU value (Y) by two and the next LRU value (Y+1) by two because DP1 just evicted location Y+1, and therefore, the new LRU value will be Y+2. This operation is illustrated at row **646**.

(511) Example methods, apparatus, systems, and articles of manufacture for eviction in a victim storage are disclosed herein. Further examples and combinations thereof include the following:

(512) Example 1 includes an apparatus comprising a cache storage, a cache controller operable to receive a first memory operation and a second memory operation concurrently, comparison logic operable to identify if the first and second memory operations missed in the cache storage, and a replacement policy component operable to, when at least one of the first and second memory operations corresponds to a miss in the cache storage, reserve an entry in the cache storage to evict based on the first and second memory operations.

(513) Example 2 includes the apparatus of example 1, wherein the replacement policy component is to speculatively lock the entry in the cache storage for eviction.

(514) Example 3 includes the apparatus of example 1, wherein the replacement policy component is operable to store an indicator of a first way of the cache storage to be replaced, in response to the first memory operation missing in the cache storage and the second memory operation hitting in the cache storage determine whether the second memory operation is directed to the first way of the

cache storage indicated by the indicator, and increment the indicator to indicate a second way of the cache storage based on the second memory operation being directed to the first way of the cache storage, and causing the second way of the cache storage to be evicted based on the first memory operation and the incremented indicator.

(515) Example 4 includes the apparatus of example 1, wherein the replacement policy component is operable to store an indicator of a second way of the cache storage to be replaced, in response to the second memory operation missing in the cache storage and the first memory operation hitting in the cache storage determine whether the first memory operation is directed to the second way of the cache storage indicated by the indicator, and decrement the indicator to indicate a first way of the cache storage based on the first memory operation being directed to the second way of the cache storage, and causing the first way of the cache storage to be evicted based on the second memory operation and the incremented indicator.

(516) Example 5 includes the apparatus of example 1, wherein the replacement policy component is operable to store a first indicator of a first way and a second indicator of a second way of the cache storage to be replaced, in response to the first memory operation missing in the cache storage and the second memory operation missing in the cache storage causing the first way of the cache storage to be evicted based on the first memory operation and the second way of the cache storage to be evicted based on the second memory operation.

(517) Example 6 includes the apparatus of example 5, wherein the replacement policy component is operable to increment the first indicator by two locations and the second indicator by two locations after the first way and the second way of the cache storage are evicted.

(518) Example 7 includes the apparatus of example 1, wherein the cache storage is a victim cache storage.

(519) Example 8 includes the apparatus of example 1, further including a first interface and a second interface, the first interface to obtain the first memory operation from a central processing unit and the second interface to obtain the second memory operation from the central processing unit, the first interface and the second interface coupled to the comparison logic and the cache controller.

(520) Example 9 includes the apparatus of example 8, wherein the first interface is a vector interface and the second interface is a scalar interface.

(521) Example 10 includes a method comprising receiving a first memory operation and a second memory operation concurrently, identifying if the first and second memory operations missed in a cache storage, and when at least one of the first and second memory operations corresponds to a miss in the cache storage, reserving an entry in the cache storage to evict based on the first and second memory operations.

(522) Example 11 includes the method of example 10, further including speculatively locking the entry in the cache storage for eviction.

(523) Example 12 includes the method of example 10, further including storing an indicator of a first way of the cache storage to be replaced, in response to the first memory operation missing in the cache storage and the second memory operation hitting in the cache storage determining whether the second memory operation is directed to the first way of the cache storage indicated by the indicator, and incrementing the indicator to indicate a second way of the cache storage based on the second memory operation being directed to the first way of the cache storage, and causing the second way of the cache storage to be evicted based on the first memory operation and the incremented indicator.

(524) Example 13 includes the method of example 10, further including storing an indicator of a second way of the cache storage to be replaced, in response to the second memory operation missing in the cache storage and the first memory operation hitting in the cache storage determining whether the first memory operation is directed to the second way of the cache storage indicated by the indicator, and decrementing the indicator to indicate a first way of the cache

storage based on the first memory operation being directed to the second way of the cache storage, and causing the first way of the cache storage to be evicted based on the second memory operation and the incremented indicator.

(525) Example 14 includes the method of example 10, further including storing a first indicator of a first way and a second indicator of a second way of the cache storage to be replaced, in response to the first memory operation missing in the cache storage and the second memory operation missing in the cache storage causing the first way of the cache storage to be evicted based on the first memory operation and the second way of the cache storage to be evicted based on the second memory operation.

(526) Example 15 includes the method of example 14, further including incrementing the first indicator by two locations and the second indicator by two locations after the first way and the second way of the cache storage are evicted.

(527) Example 16 includes a system comprising a central processing unit to concurrently output a first memory operation and a second memory operation, a cache coupled to the central processing unit, the cache further including a cache storage, a cache controller operable to receive a first memory operation and a second memory operation concurrently, comparison logic operable to identify if the first and second memory operations missed in the cache storage, and a replacement policy component operable to, when at least one of the first and second memory operations corresponds to a miss in the cache storage, reserve an entry in the cache storage to evict based on the first and second memory operations.

(528) Example 17 includes the system of example 16, wherein the cache storage is a first cache storage, the cache further including a second cache storage coupled in parallel with the first cache storage.

(529) Example 18 includes the system of example 16, wherein the cache storage is a victim cache storage.

(530) Example 19 includes the system of example 16, wherein the cache further includes a first interface and a second interface, the first interface is a 64-bit wide bidirectional scalar interface and the second interface is a 512-bit wide vector interface.

(531) Example 20 includes the system of example 16, wherein the replacement policy component is operable to adjust the entry reservations in the cache storage based on 1) a validity of the first and second memory operations, 2) whether the cache storage stores data for the first and second memory operations, and 3) whether the first and second memory operations are to allocate data to the cache storage or write data to the cache storage. FIG. 11A is an example circuit implementation of the victim cache store queue **216** of FIGS. 2 and/or 3. In FIG. 11A, the victim cache store queue **216** includes example latches **1102a**, **1102b**, **1102c**, **1102d**, **1102e**, example merge circuits **1103a-c**, an example arithmetic component **1104**, an example atomic compare component **1106**, an example read-modify-write merge component **1108**, an example select multiplexer **1110**, and example ECC generator **1112**, an example arbitration manager **1114**, an example pending store address data store **1116**, an example priority multiplexer **1118**, and an example write port **1126**. The example merge circuits **1103a-d** include an example comparator(s) **1120**, and example switches **1122**. The example of FIG. 11A illustrates a single pipeline of the victim cache store queue **216**. However, the victim storage element **216** may be arranged to support more than one independent copy of the pipeline with respect to different banks as indicated by the dashed box **1100**. Accordingly, the pipeline of FIG. 11A may be reproduced multiple times for different banks, as further described below.

(532) Methods and Apparatus for Read-Modify-Write Support in Multi-Banked Data Ram Cache for Bank Arbitration

(533) Some monolithic storage devices do not support multiple accesses by a processor (e.g., a CPU) during the same clock cycle. For example, a request to access data in a single main storage can lock up the entire single main storage. In such examples, there is a single register file capable of supporting one full cache line access per clock cycle. In some such examples, an entire cache

line associated with the single main storage can be locked to service the request because the single register file is allocated to the storage data bank that received such a request.

(534) Examples disclosed herein include multi-banked data RAM cache for bank arbitration. In some disclosed examples, the main cache store queue **212** of FIG. 2 has a plurality of banks (e.g., storage banks, storage data banks, memory banks, etc.) and, thus, is multi-banked, or has a multi-bank data or memory structure. In some disclosed examples, the main storage **214** of FIG. 2 has a plurality of banks and, thus, is multi-banked, or has a multi-bank data or memory structure. Advantageously, the multi-bank structure of the main cache store queue **212** and/or the main storage **214** can effectuate support for two or more accesses (e.g., CPU accesses) per clock cycle. Advantageously, the multi-bank structure of the main cache store queue **212** and/or the main storage **214** allows the two or more accesses to be serviced and/or fulfilled without performance overhead.

(535) In some disclosed examples, arbitration logic is coupled to the main cache storage queue **212** and/or the main storage **214** in an arrangement to resolve bank access conflicts. For example, the main cache storage queue **212** can be divided into 16 independently addressable banks. In such examples, the arbitration logic can include 16 sets of arbitration logic, where each set of the arbitration logic is in circuit with a corresponding one of the 16 independently addressable banks. Advantageously, the arbitration logic can resolve bank access conflicts of a first bank of the 16 banks independently of a second bank of the 16 banks.

(536) FIG. 7B is a schematic illustration of a second example encapsulated data cache system **710**. The second encapsulated data cache system **710** can be an example circuit implementation of the L1 cache **110** of FIG. 1 or portion(s) thereof, and/or, more generally, the data cache **108** of FIG. 1 or portion(s) thereof. The second encapsulated data cache system **710** is encapsulated to provide a unified storage view to an external system (e.g., one or more CPUs, one or more processors, external hardware, etc.). For example, the second encapsulated data cache system **710** can be encapsulated using firmware, software, hardware logic, etc., and/or a combination thereof to encapsulate and/or otherwise abstract all complexity related to read after write (RAW) hazards, write after write (WAW) hazards, data forwarding, etc. Advantageously, the unified architecture depicted in the example of FIG. 7B can provide separate ports to an external system (e.g., the CPU **102**) to simplify obtaining access to the most recent versions of data of interest.

(537) In the example of FIG. 7B, the second encapsulated data cache system **710** includes the main cache store queue **212** of FIG. 2 and the main storage **214** of FIG. 2. In FIG. 7B, the main cache store queue **212** includes a plurality of first example independently addressable banks (e.g., data banks, store banks, data store banks, etc.) **712**. For example, the first independently addressable banks **712** can be bank queues, bank sub-queues, etc. In such examples, the bank queues, the bank sub-queues, etc., can include a plurality of storage elements, as depicted below in FIG. 7D and/or FIG. 7E. The first independently addressable banks **712** are main cache store queue banks (e.g., main cache store queue data banks, main cache bank queues, main cache bank sub-queues, etc.). In FIG. 7B, the main cache store queue **212** includes 16 of the main cache store queue banks **712**. Alternatively, there may be fewer or more than 16 instances of the main cache store queue banks **712**.

(538) In FIG. 7B, the main storage **214** includes a plurality of second example independently addressable banks (e.g., data banks, store banks, data store banks, etc.) **714**. The second independently addressable banks **714** are main storage banks (e.g., main storage data banks). In FIG. 7B, the main storage **214** includes 16 of the main storage banks **714**. Alternatively, there may be fewer or more than 16 instances of the main storage banks **714**. For example, the main cache store queue **212** can include a plurality of bank queues or bank sub-queues including a first bank queue or a first bank sub-queue (e.g., MAIN CACHE STORE QUEUE: BANK 1) and the main storage **214** can include a plurality of data banks including a first data bank (e.g., MAIN STORAGE: BANK 1).

(539) In some examples, the main cache store queue **212** and the main storage **214** each have 256 rows with each row having a line width of 1024 bits, but being divided into 16 banks as illustrated in FIG. 7B, with 64 bits per row in a given bank. In some examples, the main cache store queue **212** and the main storage **214** each have 384 rows with each row having a line width of 1024 bits, but being divided into 16 banks as illustrated in FIG. 7B, with 64 bits per row in a given bank. Of course, the cache size described in the above examples are only two examples and the disclosure is not limited to any particular cache line width, number of banks, or rows, etc. As depicted in FIG. 7B, each of the first banks **712** correspond to one of the second banks **714**. For example, a first bank (BANK 1) of the main cache store queue **212** is in circuit with a first bank (BANK 1) of the main storage **214** and, thus, corresponds with the first bank of the main storage **214**.

(540) Advantageously, the multi-bank structure of the main cache store queue **212**, the main storage **214**, and/or, more generally, the second encapsulated data cache system **710**, can service read and write operations that are sent to the banks in parallel. In some examples, each bank arbitrates its own processes in response to the read and/or write operations. By allowing each bank of the second encapsulated data cache system **710** to operate independently, operation of the second encapsulated data cache system **710** is more efficient since an entire cache line is not locked up when a request is received. Rather, only the portion of the cache line allocated to the bank that received such a request would be locked.

(541) FIG. 7D depicts an example main cache multi-bank structure **730**. In some examples, the L1 cache **110**, the L2 cache **112**, and/or the L3 cache **114** of FIG. 1 can have the main cache multi-bank structure **730**. In some examples, the main cache store queue **212** of FIG. 2 and/or the main storage **214** of FIG. 2 can have the main cache multi-bank structure **730**. For example, the main cache multi-bank structure **730** can be an example implementation of the main cache store queue **212** and/or the main storage **214**.

(542) The main cache multi-bank structure **730** is a data or memory structure that includes 16 example banks (Banks 0-15) **732**, with each of the banks **732** having a data width of 64 bytes (e.g., bytes 0-7). Each of the banks **732** includes 16 example rows (Rows 0-15) **734**, with each of the rows **734** having a data width of 128 bytes (e.g., 128 bytes=8 bytes/bank*16 banks). In FIG. 7D, each of the banks **732** can store 128 bytes (e.g., 128 bytes=8 bytes/row*16 rows). In FIG. 7D, the main cache multi-bank structure **730** can store 2048 bytes (e.g., 2048 bytes=128 bytes/bank*16 banks). The main cache multi-bank structure **730** is independently addressable by bank. For example, the first row of the rows **734** has a starting row address of 0 and an ending row address of 127, a second row of the rows **734** has a starting row address of 128 and an ending row address of 255, etc.

(543) In some examples, a cache line can be 128 bytes of data that fits in a width of memory (e.g., DRAM) or storage unit (e.g., the main storage **214**, the victim storage **218**, etc.). In the example of FIG. 7D, a cache line can consume an entire row of the main cache bank structure **730**. For example, a cache line can use one of the rows **734** of 16 banks, where each bank is 8 bytes wide. Advantageously, the main cache bank structure **730** can enable 16 different cache lines to access data stored therein.

(544) FIG. 7E depicts an example unified cache multi-bank structure **740**. In some examples, the L1 cache **110**, the L2 cache **112**, and/or the L3 cache **114** of FIG. 1 can have the unified cache bank structure **740**. In some examples, the main cache store queue **212** of FIG. 2, the main storage **214** of FIG. 2, the victim cache store queue **216** of FIG. 2, and/or the victim storage **218** of FIG. 2 can have the unified cache multi-bank structure **740**. For example, the unified cache multi-bank structure **740** can be an example implementation of the main cache store queue **212**, the main storage **214**, the victim cache store queue **216**, and/or the victim storage **218**.

(545) The unified cache multi-bank structure **740** is a data or memory structure that includes 16 example banks (Banks 0-15) **742**, with each of the banks **742** having a data width of 64 bytes (e.g., bytes 0-7). Each of the banks **742** includes 16 example rows (Rows 0-15) **744**, with each of the

rows **744** having a data width of 128 bytes (e.g., 128 bytes=8 bytes/bank*16 banks). In FIG. 7E, each of the banks **742** can store 128 bytes (e.g., 128 bytes=8 bytes/row*16 rows). In FIG. 7E, the unified cache multi-bank structure **740** can store 2048 bytes (e.g., 2048 bytes=128 bytes/bank*16 banks). The unified cache multi-bank structure **740** is independently addressable by bank. For example, the first row of the rows **744** has a starting row address of 0 and an ending row address of 127, a second row of the rows **744** has a starting row address of 128 and an ending row address of 255, etc.

(546) FIGS. **9A** and **9B** (collectively FIG. **9**) depict example instruction servicing workflows **900**, **910**, **920** that can be executed by the main cache controller **222**, the victim cache controller **224**, and/or, more generally, the cache controller **220** of FIG. **2**, and/or, more generally, the L1 cache **110** of FIGS. **1** and/or **2**. For example, the instruction servicing workflows **900**, **910**, **920** can be executed in connection with at least one of the main cache store queue **212**, the main storage **214**, the victim cache store queue **216**, or the victim storage **218** of FIG. **2**. The instruction servicing workflows **900**, **910**, **920** include a first example instruction servicing workflow **900**, a second example instruction servicing workflow **910**, and a third example instruction servicing workflow **920**.

(547) The first instruction servicing workflow **900** is executed in response to receiving a first example store instruction **902**. The first store instruction **902** is a CPU write operation or command to write 64 bytes at address (Addr) 0. The address processing components **302a-c** of FIGS. **3A-3D** can rotate and/or otherwise organize first data (CPU-Write Data) included in the first store instruction **902** to generate first example rotated or aligned data (Write-Data Aligned) **904**.

(548) The address processing components **302a-c** of FIGS. **3A-3D** can use a memory address from the store instruction to determine which banks of the main cache store queue **212**, the main storage **214**, the victim cache store queue **216**, and/or the victim storage **218** of FIG. **2** are needed for the first store instruction **902**. For example, the address processing components **302a-c** can determine that Addr 0 of the first store instruction **902** is indicative of 8 entire banks (e.g., 8 of the banks **722** of FIG. **7C**) to be read from and written to. In such examples, the address processing components **302a-c** can determine that the number or quantity of banks to read from is 0x00FF and the number or quantity of banks to write to is 0x00FF. For example, the address processing components **302a-c** can determine that Banks 0-7 of FIG. **7C** need to be accessed, where each of the banks has a corresponding bit (e.g., a first bit for Bank 7, a second bit for Bank 6, a third bit for Bank 5, etc.). When a bank is to be read from or written to, the corresponding bit position has a bit value of 1 and a bit value of 0 otherwise. In such examples, the address processing components **302a-c** can generate an address for the number of banks read of 0x00FF, which is 11111111 in binary, based on each of the bits for Banks 0-7 having a 1 value (e.g., Bank 7 is 1, Bank 6 is 1, etc.) indicative of that respective bank needed to be accessed for the first store instruction **902**.

(549) In some examples, the bank processing logic **303** of FIGS. **3A-3D** detects whether incoming store instructions, such as the first store instruction **902**, indicate a write of an entire bank, or a write of a partial bank. In the first instruction servicing workflow **900**, the bank processing logic **303** can determine that, since all of the needed banks are to be completely overwritten, then none of the banks are needed to be first read from. For example, the bank processing logic **303** can determine that the number of banks read is 0x0000, which is 00000000 in binary and is indicative of each of the banks not needed to be read from. Advantageously, the bank processing logic **303** can reduce the number of banks to read from and, thus, improves efficiency and/or otherwise optimizes operation of the main cache store queue **212**, the main storage **214**, and/or, more generally, the encapsulated data cache system **700** of FIG. **7** by executing less read operations compared to previous implementations of cache systems.

(550) The second instruction servicing workflow **910** is executed in response to receiving a second example store instruction **912**. The second store instruction **912** is a CPU write operation or command to write 64 bytes at Addr 3. The address processing components **302a-c** can rotate and/or

otherwise organize second data (CPU-Write Data) included in the second store instruction **912** to generate second example rotated or aligned data (Write-Data Aligned) **914**. For example, the address processing components **302a-c** can shift the second data to place the second data in the correct lane associated with cache memory, such as the L1 cache **110**.

(551) The address processing components **302a-c** can determine that Addr 3 of the second store instruction **912** is indicative of 8 entire banks (e.g., 8 of the banks **702** of FIG. 7) to be read from and written to. In such examples, the address processing components **302a-c** can determine that the number or quantity of banks to read from is 0x01FF and the number or quantity of banks to write to is 0x01FF. For example, the address processing components **302a-c** can determine that Banks 0-8 of FIG. 7 need to be accessed, where each of the banks has a corresponding bit (e.g., a first bit for Bank 8, a second bit for Bank 7, a third bit for Bank 6, etc.). In such examples, the address processing components **302a-c** can generate an address for the number of banks read of 0x01FF, which is 11111111 in binary, based on each of the bits for Banks 0-8 having a 1 value (e.g., Bank 8 is 1, Bank 7 is 1, etc.) indicative of that respective bank needed to be accessed for the second store instruction **912**.

(552) In some examples, the bank processing logic **303** of FIGS. 3A-3D detects whether the third store instruction **922** indicates a write of an entire bank, or a write of a partial bank. In the second instruction servicing workflow **910**, the bank processing logic **303** can determine that when Banks 0 and 8 are to be completely overwritten, then Banks 0 and 8 do not need to be first read from. For example, the bank processing logic **303** can determine that the number of banks read is 0x0101, which is 100000001 in binary and is indicative of Banks 0 and 8 not needed to be read from. Advantageously, the bank processing logic **303** can reduce the number of banks to read from and, thus, improves efficiency and/or otherwise optimizes operation of the main cache store queue **212**, the main storage **214**, and/or, more generally, the encapsulated data cache system **700** of FIG. 7 by executing less read operations compared to previous implementations of cache systems.

(553) The third instruction servicing workflow **920** is executed in response to receiving a third example store instruction **922**. The third store instruction **922** is a CPU write operation or command to write 64 bytes at Addr **113**. The address processing components **302a-c** can rotate and/or otherwise organize second data (CPU-Write Data) included in the third store instruction **922** to generate third example rotated or aligned data (Write-Data Aligned) **924**. For example, the address processing components **302a-c** can shift the third data to place the third data in the correct lane associated with cache memory, such as the L1 cache **110**.

(554) The address processing components **302a-c** can determine that Addr **113** of the third store instruction **922** is indicative of 8 entire banks (e.g., 8 of the banks **722** of FIG. 7) to be read from and written to. In such examples, the address processing components **302a-c** can determine that the number or quantity of banks to read from is 0xC07F and the number or quantity of banks to write to is 0xC07F. For example, the address processing components **302a-c** can determine that Banks 0-6 and 14-15 of FIG. 7 need to be accessed, where each of the banks has a corresponding bit (e.g., a first bit for Bank 15, a second bit for Bank 14, a third bit for Bank 6, etc.). In such examples, the address processing components **302a-c** can generate an address for the number of banks read of 0xC07F, which is 1100000001111111 in binary, based on each of the bits for Banks 0-6 and 14-15 having a 1 value (e.g., Bank 15 is 1, Bank 14 is 1, Bank 6 is 1, etc.) indicative of that respective bank needed to be accessed for the third store instruction **922**.

(555) In some examples, the bank processing logic **303** of FIGS. 3A-D detects whether the third store instruction **922** indicates a write of an entire bank, or a write of a partial bank. In the third instruction servicing workflow **920**, the bank processing logic **303** can determine that Banks 6 and 14-15 are to be completely overwritten, then Banks 6 and 14-15 do not need to be first read from. For example, the bank processing logic **303** can determine that the number of banks read is 0xC080, which is 1100000001000000 in binary and is indicative of Banks 6 and 14-15 not needed to be read from. Advantageously, the bank processing logic **303** can reduce the number of banks to

read from and, thus, improves efficiency and/or otherwise optimizes operation of the main cache store queue **212**, the main storage **214**, and/or, more generally, the encapsulated data cache system **700** of FIG. 7 by executing less read operations compared to previous implementations of cache systems.

(556) FIGS. **10A-1** and **10A-2** (collectively FIG. **10A**) illustrate a schematic illustration of an example data cache system **1000**. The data cache system **1000** can be an example implementation of the L1 cache **110** of FIGS. **1**, **2**, and/or **3**, or portion(s) thereof. The data cache system **1000** includes a first example bank (ENCAPSULATED DATA CACHE SYSTEM BANK[i]) **1002** of the encapsulated data cache system **700** of FIG. 7B. For example, the first bank **1002** can correspond to MAIN CACHE STORE QUEUE: BANK 1 and MAIN STORAGE: BANK 1 of FIG. 7B. The first bank **1002** includes a first example main cache store queue bank **1004** of the main cache store queue **212** of FIG. 2, which can be an example implementation of MAIN CACHE STORE QUEUE: BANK 1 of FIG. 7B. The first bank **1002** includes a first example main storage bank **1006** of the main storage **214** of FIG. 2, which can be an example implementation of MAIN STORAGE: BANK 1 of FIG. 7B.

(557) The data cache system **1000** includes first example address processing logic **1001** and second example address processing logic **1003**. The first address processing logic **1001** and the second address processing logic **1003** can be example implementations of the address processing components **302a-c** of FIGS. 3A-3D and/or the bank processing logic **303** of FIGS. 3A-3D. The first address processing logic **1001** can be an example implementation of at least one of the first address processing component **302a**, the second address processing component **302b**, or the third address processing component **302c** of FIGS. 3A-D. The second address processing logic **1001** can be an example implementation of at least one of the first address processing component **302a**, the second address processing component **302b**, or the third address processing component **302c** of FIGS. 3A-D.

(558) The data cache system **1000** includes example arbitration logic **1008**, **1010**, and example multiplexer logic **1012**, **1014**, **1016**. The arbitration logic **1008**, **1010** includes first example arbitration logic (e.g., a first arbiter) **1008** and second example arbitration logic (e.g., a second arbiter) **1010**. The first arbitration logic **1008** is a main storage read/write arbiter (MS R/W ARB[i]) and the second arbitration logic **1010** is a main cache store queue (STQ WRITE ARB[i]). The example arbitration logic **1008**, **1010** of the illustrated example of FIG. **10A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. The example multiplexer logic **1012**, **1014**, **1016** of the illustrated example of FIG. **10A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(559) The multiplexer logic **1012**, **1014**, **1016** includes a first example multiplexer (MUX1[i]) (e.g., a first multiplexer logic circuit) **1012**, a second example multiplexer (MUX2[i]) (e.g., a second multiplexer logic circuit) **1014**, and a third example multiplexer (MUX3[i]) (e.g., a third multiplexer logic circuit) **1016**. The multiplexers **1012**, **1014**, **1016** have a select input (SEL[i]), data inputs (1-3), and an output.

(560) The first data inputs (data input 1) of the multiplexers **1012**, **1014**, **1016** are coupled to the first address processing logic **1001** and the second address processing logic **1003**. The second data inputs (data input 2) of the multiplexers **1012**, **1014**, **1016** are coupled to the first address processing logic **1001** and the second address processing logic **1003**. The third data inputs (data

input 3) of the multiplexers **1012**, **1014**, **1016** are coupled to the first address processing logic **1001** and the second address processing logic **1003**.

(561) The select input of the first multiplexer **1012** is coupled to an output of the second arbitration logic **1010**. The select input of the second multiplexer **1014** and the select input of the third multiplexer **1016** are coupled to outputs of the first arbitration logic **1008**. The output of the first multiplexer **1012** is coupled to an example write port (WRITE PORT[i]) **1024** of the first main cache store bank **1004**. The output of the second multiplexer **1014** is coupled to an example read port (READ PORT[i]) **1026** of the first main cache store bank **1004**. The output of the third multiplexer **1016** is coupled to an example read/write port (READ/WRITE PORT[i]) **1028** of the first main storage bank **1006**. The first arbitration logic **1008** is coupled to the first address processing logic **1001**, the second address processing logic **1003**, the second arbitration logic **1010**, and outputs of the first main cache store queue bank **1004**.

(562) STQ[i] of FIG. **10A** is representative of a single bank of a multi-bank implementation of the main cache store queue **212**. For example, the main cache store queue **212** can have STQ[0]-STQ[15] representative of the main cache store queue **212** having 16 banks. In such examples, each of STQ[0]-STQ[15] can store 64 bits (i.e., 8 bytes). In some such examples, STQ[0]-STQ[15], and/or, more generally, the main cache store queue **212**, can store 24,576 bits (i.e., 3072 bytes). Alternatively, each of STQ[0]-STQ[15] may store a different quantity of bits and, thus, the main cache store queue may store a different quantity of bits.

(563) A plurality of the banks including the first bank **1002**, can be encapsulated to form and/or otherwise generate an encapsulated data cache system **1034**. For example, the encapsulated data cache system **1034** can be an example implementation of the encapsulated data cache system **700** of FIG. **7B**. Advantageously, each corresponding bank of the main cache store queue **212** and the main storage **214** can be encapsulated together to form and/or otherwise generate example encapsulated data cache banks **1036** for simplification when interacting with external system(s). Each of the encapsulated data cache banks **1036** include an example encapsulated write port (WRITE PORT (STQ[i])) **438**, an example encapsulated read port (READ PORT (STQ[i])) **1040**, and an example encapsulated read/write port (READ/WRITE PORT MS[i])) **1042**.

(564) The first address processing logic **1001** and/or the second address processing logic **1003** can obtain example store instructions **1018**, **1020**, **1022** from one(s) of the interface(s) of FIG. **1** of the CPU **102** of FIG. **1**. The store instructions **1018**, **1020**, **1022** include example data, such as WDATA, ADDR, BYTEN, SIZE, and R/W. WDATA corresponds to data (e.g., 64 bits of data) to be written and/or otherwise stored in at least one of the main cache store queue **212** or the main storage **214**. ADDR corresponds to a data address associated with at least one of the main cache store queue **212** or the main storage **214**. BYTEN corresponds to byte enable data. SIZE corresponds to a data size of a data access operation (e.g., a read operation, a write operation, a modify operation, etc., and/or a combination thereof). R/W corresponds to whether the store instruction is a read operation or a write operation.

(565) The store instructions **1018**, **1020**, **1022** include a first example store instruction (SCALAR_DP (DP0)) **1018**, a second example store instruction (DMA) **1020**, and a third example store instruction (VECTOR_DP (DP1)) **1022**. The first store instruction **1018** is transmitted from the scalar interface of FIG. **1** (e.g., the CPU interface **202** of FIG. **2**) and, thus, corresponds to a scalar data path (SCALAR_DP (DP1)) of the data cache system **1000**. The second store instruction **1020** is transmitted from the memory interface of FIG. **1**, such as a direct memory access (DMA) interface and, thus, corresponds to a DMA data path (DMA). The third store instruction **1022** is transmitted from the vector interface of FIG. **1** (e.g., the CPU interface **202** of FIG. **2**) and, thus, corresponds to a vector data path (VECTOR_DP (DP1)) of the data cache system **1000**.

(566) The first address processing logic **1001** and/or the second address processing logic **1003** generate transaction data (TRANSACTION_DP0[i], TRANSACTION DMA[i], TRANSACTION DP1[i]) that can be used to execute a data access operation associated with at least one of the main

cache store queue **212** or the main storage **214**. The first address processing logic **1001** can extract, and in some examples rotate, the WDATA from respective one(s) of the store instructions **1018**, **1020**, **1022** and transmit the extracted and/or rotated WDATA to a respective first input of the multiplexers **1012**, **1014**, **1016**. For example, the first address processing logic **1001** can extract and rotate first WDATA from the first store instruction **1018** and transmit the first extracted and rotated WDATA to the first input of the first multiplexer **1012**, the first input of the second multiplexer **1014**, and the first input of the third multiplexer **1016**.

(567) The second address processing logic **1003** can determine an address (MS/STQ_ADDR[i]) for one or more of the 16 banks of at least one of the store queue **212** or the main storage **214**. The address can be based on the ADDR data included in the store instructions **1018**, **1020**, **1022**. The second address processing logic **1003** can determine a byte enable value per bank (BYTEN/BANK[i]) based on the BYTEN data included in the store instructions **1018**, **1020**, **1022**. The second address processing logic **1003** can determine a write bank request (WR_BANK_REQ[i]) and/or a read bank request (RD_BANK_REQ[i]) based on the R/W data included in the store instructions **1018**, **1020**, **1022**.

(568) The first address processing logic **1001** and/or the second address processing logic **1003** can determine transaction data for respective ones of the store instructions **1018**, **1020**, **1022**. The transaction data can include the rotated WDATA data, MS/STQ_ADDR[i], and BYTEN/BANK[i]. For example, the first address processing logic **1001** and/or the second address processing logic **1003** can generate first transaction data (TRANSACTION_DP0[i]) based on the first store instruction **1018**, second transaction data (TRANSACTION_DMA[i]) based on the second store instruction **1020**, and third transaction data (TRANSACTION_DP1[i]) based on the third store instruction **1022**. In such examples, the first address processing logic **1001** and/or the second address processing logic **1003** can transmit the first transaction data to the first inputs of the multiplexers **1012**, **1014**, **1016**, the second transaction data to the second inputs of the multiplexers **1012**, **1014**, **1016**, and the third transaction data to the third inputs of the multiplexers **1012**, **1014**, **1016**.

(569) In example operating conditions, the first address processing logic **1001** and the second address processing logic **1003** obtain the store instructions **1018**, **1020**, **1022**. The first address processing logic **1001** and the second address processing logic **1003** generate the first through third transaction data based on respective ones of the store instructions **1018**, **1020**, **1022**. The first address processing logic **1001** and the second address processing logic **1003** transmit the first through third transaction data to the multiplexers **1012**, **1014**, **1016**. The second address processing logic **1003** transmit either a read bank request or a write bank request corresponding to each of the store instructions **1018**, **1020**, **1022**.

(570) In example operating conditions, the first arbitration logic **1008** determines whether one(s) of the store instructions **1018**, **1020**, **1022** are requesting to read one or more banks of the main cache store queue **212** or write to one or more banks of the main storage **214**. In example operating conditions, the first arbitration logic **1008** prioritizes read operations over write operations. In example operating conditions, the bank(s) of the store queue **212** can generate an example store queue full signal (FULL_SIG[i]) **1030** in response to the store queue **212** being full. In example operating conditions, the bank(s) of the store queue **212** can generate an example complete data write signal (COMPLETE_DATA_WR_SIG[i]) **1032**.

(571) For example, the first store instruction **1018** can correspond to a write operation for Banks 0-4, the second store instruction **1020** can correspond to a read operation of Banks 5-9, and the third store instruction **1022** can correspond to a read operation of Banks 10-14. The second arbitration logic **1010** can assign DP0 to transmit the first transaction data to the write port **1024** of Banks 0-4 (e.g., WRITE PORT[0], WRITE PORT[1], WRITE PORT[2], etc.) because no other data paths are requesting a write operation to be serviced. The second arbitration logic **1010** can assign DP0 by generating a signal (SEL[i]) to instruct the first multiplexer **1012** to select the first transaction data.

The first arbitration logic **1008** can assign DMA to transmit the second transaction data to the read port **1026** of Banks 5-9 (e.g., READ PORT[4], READ PORT[5], READ PORT[6], etc.) because no other data paths are requesting a read operation to be serviced in connection with Banks 5-9. The first arbitration logic **1008** can assign DMA by generating a signal (SEL[i]) to instruct the second multiplexer **1014** to select the second transaction data. The first arbitration logic **1008** can assign DP1 to transmit the third transaction data to the read port **1026** of Banks 10-14 (e.g., READ PORT[9], READ PORT[10], READ PORT[11], etc.) because no other data paths are requesting a read operation to be serviced in connection with Banks 10-14. The first arbitration logic **1008** can assign DP1 by generating a signal (SEL[i]) to instruct the third multiplexer **1016** to select the third transaction data. In such examples, Bank 15 is not assigned to any of the data paths.

(572) In example operating conditions, if the first store instruction **1018** corresponds to a write operation for Banks 0-4 that is split, the first arbitration logic **1008** can delay and/or otherwise stall the write operation. For example, if a first portion of the write operation is associated with writing to the main cache store queue **212** and a second portion of the write operation is associated with reading from the main storage **214**, the first arbitration logic **1008** can instruct the second arbitration logic **1010** to not service and/or otherwise not assign the first transaction data to the write port **1024**.

(573) In example operating conditions, if the first store instruction **1018** corresponds to a write operation for Banks 0-4 that is split, the main cache store queue **212** can instruct the first arbitration logic **408** to service the write operation when the complete data has been assembled for writing. For example, if a first portion of the write operation is associated with writing to the main cache store queue **212** and a second portion of the write operation is associated with reading from at least one of the main cache store queue **212** or the main storage **214**, the first arbitration logic **1008** can wait to assign the first transaction data to the read/write port **1028**. In such examples, in response to locating data associated with the second portion in the main cache store queue **212**, the main cache store queue **212** can deliver the located data to the main storage **214**. In response to delivering the data to the main storage **214**, the main cache store queue **212** can generate a signal (e.g., assert a logic high signal) for COMPLETE_DATA_WR_SIG[i] instructing the first arbitration logic **1010** to service the write operation because the complete set of data required for the write operation has been read and/or otherwise assembled for servicing.

(574) FIGS. **10B-1** and **10B-2** (collectively FIG. **10B**) illustrate a schematic illustration of an example data cache system **1000b**. The data cache system **1000b** can be an example implementation of the L1 cache **110** of FIGS. **1**, **2**, and/or **3**, or portion(s) thereof. The data cache system **1000b** includes a first example bank (ENCAPSULATED DATA CACHE SYSTEM BANK[i]) **1002b** of the encapsulated data cache system **700** of FIG. **7A**. For example, the first bank **1002b** can correspond to VICTIM CACHE STORE QUEUE: BANK 1 and VICTIM STORAGE: BANK 1 of FIG. **7A**. The first bank **1002b** includes a first example victim cache store queue bank **1004b** of the victim cache store queue **216** of FIG. **2**, which can be an example implementation of VICTIM CACHE STORE QUEUE: BANK 1 of FIG. **7A**. The first bank **1002b** includes a first example victim storage bank **1006b** of the victim storage **218** of FIG. **2**, which can be an example implementation of VICTIM STORAGE: BANK 1 of FIG. **7A**.

(575) The data cache system **1000b** includes first example address processing logic **1001b**, second example address processing logic **1003b**, example arbitration logic **1008b**, **1010b**, and example multiplexer logic **1012b**, **1014b**, **1016b**. The arbitration logic **1008b**, **1010b** includes first example arbitration logic (e.g., a first arbiter) **1008b** and second example arbitration logic (e.g., a second arbiter) **1010b**. The first arbitration logic **1008b** is a victim storage read/write arbiter (VS R/W ARB[i]) and the second arbitration logic **1010b** is a victim cache store queue (STQ_V WRITE ARB [i]). The example arbitration logic **1008b**, **1010b** of the illustrated example of FIG. **10B** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog

or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. The example multiplexer logic **1012b**, **1014b**, **1016b** of the illustrated example of FIG. **10B** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(576) The multiplexer logic **1012b**, **1014b**, **1016b** includes a first example multiplexer (MUX1[i]) (e.g., a first multiplexer logic circuit) **1012b**, a second example multiplexer (MUX2[i]) (e.g., a second multiplexer logic circuit) **1014b**, and a third example multiplexer (MUX3[i]) (e.g., a third multiplexer logic circuit) **1016b**. The multiplexers **1012b**, **1014b**, **1016b** have a select input (SEL[i]), data inputs (1-3), and an output.

(577) The first data inputs (data input 1) of the multiplexers **1012b**, **1014b**, **1016b** are coupled to the address processing logic **1001b**, **1003b**. The second data inputs (data input 2) of the multiplexers **1012b**, **1014b**, **1016b** are coupled to the address processing logic **1001b**, **1003b**. The third data inputs (data input 3) of the multiplexers **1012b**, **1014b**, **1016b** are coupled to the address processing logic **1001b**, **1003b**.

(578) The select input of the first multiplexer **1012b** is coupled to an output of the second arbitration logic **1010b**. The select input of the second multiplexer **1014b** and the select input of the third multiplexer **1016b** are coupled to outputs of the first arbitration logic **1008b**. The output of the first multiplexer **1012b** is coupled to an example write port (WRITE PORT[i]) **1024b** of the first victim cache store bank **1004b**. The output of the second multiplexer **1014b** is coupled to an example read port (READ PORT[i]) **1026b** of the first victim cache store bank **1004b**. The output of the third multiplexer **1016b** is coupled to an example read/write port (READ/WRITE PORT[i]) **1028b** of the first victim storage bank **1006b**. The first arbitration logic **1008b** is coupled to the address processing logic **1001b**, **1003b**, the second arbitration logic **1010b**, and outputs of the first victim cache store queue bank **1004b**.

(579) STQ_V[i] of FIG. **10B** is representative of a single bank of a multi-bank implementation of the victim cache store queue **216**. For example, the victim cache store queue **216** can have STQ_V[0]-STQ_V[15] representative of the victim cache store queue **216** having 16 banks. In such examples, each of STQ_V[0]-STQ_V[15] can store 64 bits (i.e., 8 bytes). In some such examples, STQ_V[0]-STQ_V[15], and/or, more generally, the victim cache store queue **216**, can store 24,576 bits (i.e., 3072 bytes). Alternatively, each of STQ_V[0]-STQ_V[15] may store a different quantity of bits and, thus, the victim cache store queue **216** may store a different quantity of bits.

(580) A plurality of the banks including the first bank **1002b**, can be encapsulated to form and/or otherwise generate an encapsulated data cache system **1034b**. For example, the encapsulated data cache system **1034b** can be an example implementation of the encapsulated data cache system **700** of FIG. **7A**. Advantageously, each corresponding bank of the victim cache store queue **216** and the victim storage **218** can be encapsulated together to form and/or otherwise generate example encapsulated data cache banks **1036b** for simplification when interacting with external system(s). Each of the encapsulated data cache banks **1036b** include an example encapsulated write port (WRITE PORT (STQ_V[i])) **1038b**, an example encapsulated read port (READ PORT (STQ_V[i])) **1040b**, and an example encapsulated read/write port (READ/WRITE PORT MS[i])) **1042b**.

(581) The address processing logic **1001b**, **1003b** can obtain example store instructions **1018b**, **1020b**, **1022b** from one(s) of the interface(s) of FIG. **1** of the CPU **102** of FIG. **1**. The store instructions **1018b**, **1020b**, **1022b** include example data, such as WDATA, ADDR, BYTEN, SIZE, AND R/W. WDATA corresponds to data (e.g., 64 bits of data) to be written and/or otherwise stored

in at least one of the victim cache store queue **216** or the victim storage **218**. ADDR corresponds to a data address associated with at least one of the victim cache store queue **216** or the victim storage **218**. BYTEN corresponds to byte enable data. SIZE corresponds to a data size of a data access operation (e.g., a read operation, a write operation, a modify operation, etc., and/or a combination thereof). R/W corresponds to whether the store instruction is a read operation or a write operation. (582) The store instructions **1018b**, **1020b**, **1022b** include a first example store instruction (SCALAR_DP (DP0)) **1018b**, a second example store instruction (DMA) **1020b**, and a third example store instruction (VECTOR_DP (DP1)) **1022b**. The first store instruction **1018b** is transmitted from the scalar interface of FIG. 1 (e.g., the CPU interface **202** of FIG. 2) and, thus, corresponds to a scalar data path (SCALAR_DP (DP1)) of the data cache system **1000b**. The second store instruction **1020b** is transmitted from the memory interface of FIG. 1, such as a direct memory access (DMA) interface and, thus, corresponds to a DMA data path (DMA). The third store instruction **1022b** is transmitted from the vector interface of FIG. 1 (e.g., the CPU interface **202** of FIG. 2) and, thus, corresponds to a vector data path (VECTOR_DP (DP1)) of the data cache system **1000b**.

(583) The address processing logic **1001b**, **1003b** can generate transaction data (TRANSACTION_DP0[i], TRANSACTION DMA[i], TRANSACTION DP1[i]) that can be used to execute a data access operation associated with at least one of the victim cache store queue **216** or the victim storage **218**. The first address processing logic **1001b** can extract, and in some examples rotate, the WDATA from respective one(s) of the store instructions **1018b**, **1020b**, **1022b** and transmit the extracted and/or rotated WDATA to a respective first input of the multiplexers **1012b**, **1014b**, **1016b**. For example, the first address processing logic **1001b** can extract and rotate first WDATA from the first store instruction **1018b** and transmit the first extracted and rotated WDATA to the first input of the first multiplexer **1012b**, the first input of the second multiplexer **1014b**, and the first input of the third multiplexer **1016b**.

(584) The second address processing logic **1003b** can determine an address (MS/STQ_V_ADDR[i]) for one or more of the 16 banks of at least one of the victim cache store queue **216** or the victim storage **218**. The address can be based on the ADDR data included in the store instructions **1018b**, **1020b**, **1022b**. The second address processing logic **1003b** can determine a byte enable value per bank (BYTEN/BANK[i]) based on the BYTEN data included in the store instructions **1018b**, **1020b**, **1022b**. The second address processing logic **1003b** can determine a write bank request (WR_BANK_REQ[i]) and/or a read bank request (RD_BANK_REQ[i]) based on the R/W data included in the store instructions **1018b**, **1020b**, **1022b**.

(585) The address processing logic **1001b**, **1003b** can determine transaction data for respective ones of the store instructions **1018b**, **1020b**, **1022b**. The transaction data can include the rotated WDATA data, MS/STQ_V_ADDR[i], and BYTEN/BANK[i]. For example, the address processing logic **1001b**, **1003b** can generate first transaction data (TRANSACTION_DP0[i]) based on the first store instruction **1018b**, second transaction data (TRANSACTION_DMA[i]) based on the second store instruction **1020b**, and third transaction data (TRANSACTION_DP1[i]) based on the third store instruction **1022b**. In such examples, the address processing logic **1001b**, **1003b** can transmit the first transaction data to the first inputs of the multiplexers **1012b**, **1014b**, **1016b**, the second transaction data to the second inputs of the multiplexers **1012b**, **1014b**, **1016b**, and the third transaction data to the third inputs of the multiplexers **1012b**, **1014b**, **1016b**.

(586) In example operating conditions, the address processing logic **1001b**, **1003b** obtain the store instructions **1018b**, **1020b**, **1022b**. The address processing logic **1001b**, **1003b** generate the first through third transaction data based on respective ones of the store instructions **1018b**, **1020b**, **1022b**. The address processing logic **1001b**, **1003b** transmit the first through third transaction data to the multiplexers **1012b**, **1014b**, **1016b**. The second address processing logic **1003b** transmit either a read bank request or a write bank request corresponding to each of the store instructions **1018b**, **1020b**, **1022b**.

(587) In example operating conditions, the first arbitration logic **1008b** determines whether one(s) of the store instructions **1018b**, **1020b**, **1022b** are requesting to read one or more banks of the victim cache store queue **216** or write to one or more banks of the victim storage **218**. In example operating conditions, the first arbitration logic **1008b** prioritizes read operations over write operations. In example operating conditions, the bank(s) of the victim cache store queue **216** can generate an example store queue full signal (FULL_SIG[i]) **1030b** in response to the victim cache store queue **216** being full. In example operating conditions, the bank(s) of the victim cache store queue **216** can generate an example complete data write signal (COMPLETE_DATA_WR_SIG[i]) **1032b**.

(588) For example, the first store instruction **1018b** can correspond to a write operation for Banks 0-4, the second store instruction **1020b** can correspond to a read operation of Banks 5-9, and the third store instruction **1022b** can correspond to a read operation of Banks 10-14. The second arbitration logic **1010b** can assign DP0 to transmit the first transaction data to the write port **1024b** of Banks 0-4 (e.g., WRITE PORT[0], WRITE PORT[1], WRITE PORT[2], etc.) because no other data paths are requesting a write operation to be serviced. The second arbitration logic **1010b** can assign DP0 by generating a signal (SEL[i]) to instruct the first multiplexer **1012b** to select the first transaction data. The first arbitration logic **1008b** can assign DMA to transmit the second transaction data to the read port **1026b** of Banks 5-9 (e.g., READ PORT[4], READ PORT[5], READ PORT[6], etc.) because no other data paths are requesting a read operation to be serviced in connection with Banks 5-9. The first arbitration logic **1008b** can assign DMA by generating a signal (SEL[i]) to instruct the second multiplexer **1014b** to select the second transaction data. The first arbitration logic **1008b** can assign DP1 to transmit the third transaction data to the read port **1026b** of Banks 10-14 (e.g., READ PORT[9], READ PORT[10], READ PORT[11], etc.) because no other data paths are requesting a read operation to be serviced in connection with Banks 10-14. The first arbitration logic **1008b** can assign DP1 by generating a signal (SEL[i]) to instruct the third multiplexer **1016b** to select the third transaction data. In such examples, Bank 15 is not assigned to any of the data paths.

(589) In example operating conditions, if the first store instruction **1018b** corresponds to a write operation for Banks 0-4 that is split, the first arbitration logic **1008b** can delay and/or otherwise stall the write operation. For example, if a first portion of the write operation is associated with writing to the victim cache store queue **216** and a second portion of the write operation is associated with reading from the victim storage **218**, the first arbitration logic **1008b** can instruct the second arbitration logic **1010b** to not service and/or otherwise not assign the first transaction data to the write port **1024b**.

(590) In example operating conditions, if the first store instruction **1018b** corresponds to a write operation for Banks 0-4 that is split, the victim cache store queue **216** can instruct the first arbitration logic **1008b** to service the write operation when the complete data has been assembled for writing. For example, if a first portion of the write operation is associated with writing to the victim cache store queue **216** and a second portion of the write operation is associated with reading from at least one of the victim cache store queue **216** or the victim storage **218**, the first arbitration logic **1008b** can wait to assign the first transaction data to the read/write port **1028b**. In such examples, in response to locating data associated with the second portion in the victim cache store queue **216**, the victim cache store queue **216** can deliver the located data to the victim storage **218**. In response to delivering the data to the victim storage **218**, the victim cache store queue **216** can generate a signal (e.g., assert a logic high signal) for COMPLETE_DATA_WR_SIG[i] instructing the first arbitration logic **1010b** to service the write operation because the complete set of data required for the write operation has been read and/or otherwise assembled for servicing.

(591) In FIG. 11A, the example latches **1102a**, **1102b**, **1102c**, **1102d**, **1102e**, are electronic devices configured to store information (e.g., bytes, bits, etc.) obtained by the victim cache store queue **216**. The example latches **1102a-c** pass the write data and information corresponding to whether the

write data needs to be combined with the read and corrected data out of the ECC logic **310** in arithmetic unit **404**, the atomic unit **406**, and/or the RMW merge component **408**. In the example of FIG. **11A**, the latch **1102a** is communicatively coupled to the cache controller **220** to obtain read, write, and/or modify instructions. Such read, modify, and/or write instructions may be originate from the CPU **102**, and transmitted to the latch **1102a** via the cache controller **220**. Latch **1102a** is coupled to latch **1102b**, the tag RAM **210**, the arbitration manager **1114**, and the pending store address data store **1116** to transmit such read, modify, and/or write instructions to the latch **1102b**, the tag RAM **210**, the arbitration manager **1114**, and the pending store address data store **1116** in response to a subsequent clock cycle of the cache controller **220**.

(592) In the example of FIG. **11A**, the latch **1102b** is coupled to latch **1102a**, the pending store address data store **1116**, latch **1102c**, the priority multiplexer **1118**, and the comparator(s) **1120** to transmit the obtained read, modify, and/or write instruction in response to a subsequent clock cycle of the cache controller **220**. In this manner, the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.), propagates through the victim cache store queue **216**.

(593) The example latch **1102c** is coupled to the latch **1102b**, the priority multiplexer **1118**, the arithmetic component **1104**, the atomic compare component **1106**, and the read-modify-write merge component **1108**. This coupling enables the latch **1102c** to transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to the arithmetic component **1104**, the atomic compare component **1106**, and/or the read-modify-write merge component **1108** in response to a subsequent clock cycle of the cache controller **220**. In some examples disclosed herein, latch **1102c** may transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to one or more of the arithmetic component **1104**, the atomic compare component **1106**, and/or the read-modify-write merge component **1108**. For example, in the event the cache controller **220** transmits an instruction to write a bit into a currently stored word, the latch **1102c** may transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to the read-modify-write merge component **1108**. In other examples disclosed herein, though the cache controller **220** may transmit an instruction to write a bit into a currently stored word, the latch **1102c** may transmit the value obtained from the read, modify, and/or write instruction (e.g., the byte value, the bit value, etc.) to the arithmetic component **1104**, the atomic compare component **1106**, and/or read-modify-write merge component **1108**.

(594) The example latch **1102d** is coupled to the first multiplexer **1110**, the ECC generator **1112**, and the latch **1102e**. This coupling enables the latch **1102d** to transmit the value obtained from first multiplexer **1110** to ECC generator **1112** and/or the latch **1102f** in response to a subsequent clock cycle of the cache controller **220**.

(595) The example latch **1102e** is coupled to priority multiplexer **1118** and to the MUX circuit **316** of FIGS. **3** and/or **11A**. This coupling enables the latch **1102e** to transmit the value obtained from the from the priority multiplexer **1118** to MUX circuit **316** in response to a subsequent clock cycle of the cache controller **220**. For example, on a read operation by the CPU **102** to an address that hits the victim storage element **218**, a previous write to the same address could still be inflight (e.g., not completely written to the victim storage element **218**). In such an example, the latch **1102e** collects the unwritten data byte(s) and provides them to the MUX **316** which eventually get sent back to the CPU **102**.

(596) The example merging circuit **1103a** is coupled to the latch **1102d**, the merging circuit **1103b**, the arithmetic component **1104**, the atomic compare component **1106**, and the read-modify-write merge component **1108**. The example merging circuit **1103b** is coupled to the merging circuit **1103a**, the priority multiplexer **1118**, and the merging circuit **1103c**. The example merging circuit **1103c** is coupled to the merging circuit **1103b** and the latch **1102b**. The example merging circuits **1103a-c** facilitate the comparison of read operations in different sections of the victim cache store

queue **216** to potentially reroute write operations to be merged with write operations corresponding to the same memory address location, as further described below. Although the example of FIG. **11A** includes three merging circuits **1103a-c**, there may be additional merging circuits to merge write operations from other sections of the victim cache store queue **216** (e.g., a merging circuit coupling the output of the latch **1102d** to the output of latch **1102b** and/or latch **1102a**, etc.). In some examples, the merging circuits **1103a-c** is combined into a single circuit that compares the write operations from the different latches **1102b-d** and reroutes based on matching memory addresses in any two or more of the different latches **1102b-d**.

(597) In the example illustrated in FIG. **11A**, the arithmetic component **1104** is coupled to the latch **1102c**, the first multiplexer **1110**, and to the ECC logic **312** to perform arithmetic operations on (e.g., increment, decrement, etc.) data from the victim storage **218**. Additionally, the arithmetic component **1104** performs histogram operations on the data stored in the victim storage **218**. The example arithmetic component **1104** of the illustrated example of FIG. **11A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. Operation of the example arithmetic component **1104** is further described below.

(598) In the example illustrated in FIG. **11A**, the atomic compare component **1106** is coupled to the latch **1102c**, the first multiplexer **1110**, and to the ECC logic **312** to compare data at a memory address to a key and, in the event the data at the memory address matches the key, replace the data. The example atomic compare component **1106** the illustrated example of FIG. **11A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. Operation of the example atomic compare component **1106** is further described below.

(599) In the example illustrated in FIG. **11A**, the read-modify-write merge component **1108** is coupled to the latch **1102c**, the first multiplexer **1110**, and to the ECC logic **312** to facilitate the read, modify, and/or write instruction(s) sent by the cache controller **220**. For example, the read-modify-write merge component **1108** is coupled to the ECC logic **312** to obtain the currently stored word that is to be affected by the read, modify, and/or write instruction(s). In an example write operation, the read-modify-write merge component **1108** is configured to update the currently stored word obtained from the ECC logic **312** with the new bit(s), byte(s), etc., obtained from the latch **1102c**. Additional description of the read-modify-write merge component **1108** is described below. The example read-modify-write merge component **1108** of the illustrated example of FIG. **11A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(600) The example first multiplexer **1110** is coupled to the arithmetic component **1104**, the atomic compare component **1106**, and the read-modify-write merge component **1108** to transmit, based on an indication from the cache controller **220**, the output of either the arithmetic component **1104**, the atomic compare component **1106**, or the read-modify-write merge component **1108** to the latch **1102d**. For example, in the event the cache controller **220** indicates to perform a write function (e.g., the cache control transmits a write request to the latch **1102b**), an indication is sent by the cache controller **220** to the first multiplexer **1110** to select the input connected to the read-modify-

write merge component **1108** to be transmitted to the latch **1102d**. The example first multiplexer **1110** of the illustrated example of FIG. **11A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc. (601) In the example illustrated in FIG. **11A**, the ECC generator **1112** is coupled to the latch **1102d** and to the latch **1102f** to facilitate error detection and correction in the value (e.g., byte(s), bit(s), etc.) stored in the latch **1102d**. For example, the ECC generator **1112** is configured to regenerate the ECC value (e.g., generate error detection code) which will be stored with the data (e.g., the merged word output from the read-modify-write merge component **1108**). The ECC value is used by the error detection and correction circuit to determine whether the error occurred during a read and/or write operation, as further described above. The example ECC generator **1112** of the illustrated example of FIG. **11A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(602) In FIG. **11A**, the example arbitration manager **1114** is coupled to the latch **1102a**, the latch **1102b**, the pending store address datastore **1116**, and the victim storage **218** to facilitate the read, modify, and/or write instructions obtained from the cache controller **220**. For example, in the event a write instruction is received from the cache controller **220**, the arbitration manager **1114** is configured to transmit a read instruction of the corresponding currently stored word to the victim storage **218**. In addition, the arbitration manager **1114** is coupled to the victim storage **218** to arbitrate between conflicting accesses of the victim storage **218**. When multiple operations attempt to access the victim storage **218** in the same cycle, the arbitration manager **1114** may select which operation(s) are permitted to access the victim storage **218** according to a priority scheme. Suitable priority schemes are described in more detail in below; however, in one example, the arbitration prioritizes read operations over write operations because write data that is in the victim cache store queue **216** is available for use by subsequent operations even before it is written to the victim storage **218**. Thus, there is minimal performance impact in allowing the write data to wait in the victim cache store queue **216**. However, as the victim cache store queue **216** fills with write data that has not yet been written back, the priority of the write operations may increase until they are prioritized over competing read operations.

(603) The example arbitration manager **1114** of the illustrated example of FIG. **11A** is implemented by a logic circuit such as, for example, a hardware processor. However, any other type of circuitry may additionally or alternatively be used such as, for example, one or more analog or digital circuit(s), logic circuits, programmable processor(s), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)), field programmable logic device(s) (FPLD(s)), digital signal processor(s) (DSP(s)), etc.

(604) In the example of FIG. **11A**, the pending store address data store **1116** is configured to store the address of the read, modify, and/or write instruction obtained from the cache controller **220**. In this manner, the pending store address datastore **1116** maintains a log of the addresses associated with each value stored in any of the latches **1102a**, **1102b**, **1102c**, **1102d**, **1102e**, **1102f**, **1102g**, and/or **1102h**. The example pending store address datastore **1116** of the illustrated example of FIG. **11A** may be implemented by any device for storing data such as, for example, flash memory, magnetic media, optical media, etc. Furthermore, the data stored in the pending store address datastore **1116** may be in any data format such as, for example, binary data, comma delimited data, tab delimited data, structured query language (SQL) structures, etc.

(605) The example priority multiplexer **1118** is coupled to the latch **1102b**, the latch **1102c**, the

latch **1102d**, and the latch **1102f** to facilitate read operations in the event either of the of the latch **1102b**, the latch **1102c**, the latch **1102d**, or the latch **1102f** are storing a value corresponding to a write instruction. For example, the cache controller may initiate the following four write instructions regarding a four bit word having addresses A3, A2, A1, and A0: write address A0 with the byte 0x11, write address A1 with the byte 0x22, write address A3 with the byte 0x23, and write address A0 with the byte 0x44. In such an example, in the event the cache controller **220** initiates a read instruction to read the entire four bit word, the priority multiplexer **1118** is configured to obtain the byte value 0x11 stored in the latch **1102f**, the byte value 0x22 stored in the latch **1102d**, the byte value 0x23 stored in the latch **1102c**, and the byte value 0x22 stored in the latch **1102b**. Additionally, the pending store address data store **1116** transmits an instruction to the priority multiplexer **1118** indicating which address value in is associated with the byte value stored in the latch **1102b**, the latch **1102c**, the latch **1102d**, and the latch **1102f**. In such an example, since the byte value stored in the latch **1102b** is associated with the address A0 and the byte value stored in the latch **1102f** is also associated with the address A0, the priority multiplexer **1118** is configured to transmit a packet to the latch **1102e** indicating that address A0 is 0x44 (e.g., the most recent write instruction associated with the address A0), address A1 is 0x22, and address A3 is 0x23.

(606) In this manner, the MUX circuit **316** is configured to update the value of the currently stored word with the byte values obtained from the priority multiplexer **1118**. Such an operation ensures that a read instruction transmitted by the victim cache store queue **216** probably indicates the correct word, even though the write instructions may not have fully propagated through the victim cache store queue **216**.

(607) An example read path (e.g., the read input to the tag RAM **210**) may run in parallel with the victim cache store queue **216**. Because a read operation (e.g., a read instruction) may refer to data in a write operation (e.g., a write instruction) that may not have completed yet, the victim cache store queue **216** may include write forwarding functionality that allows the read path to obtain data from the victim cache store queue **216** that has not yet been written back to the victim storage **218**. In an example, the victim cache store queue **216** includes pending store address data store **1116** that records the addresses of the operations at each stage of the victim cache store queue **216**, a priority multiplexer **1118** to select data from one of the stages (e.g., latches) of the victim cache store queue **216** for forwarding, and a MUX circuit **316** that selects between the output of the victim storage **218** (by way of the error detection and correction circuit **312**) and the forwarded victim cache store queue **216** data from the data priority multiplexer **1118**.

(608) The example write port **1126** is coupled to the write path and the latch **1102a**. The write port **1126** may be implemented by an interface that interfaces with the victim cache controller **224** (e.g., the cache controller **220**) to obtain a write instruction. For example, the write port **1126** is utilized to receive addresses and values from the cache controller **220** to write.

(609) An example flow of a read operation through the victim cache store queue **216** will now be described. In a first cycle, the L1 data cache **110** retrieves a record from the tag RAM **210** that is associated with an address of the read operation to determine whether the data is stored in the victim storage **218**. In a direct mapped example, the L1 data cache **110** need not wait for the tag RAM **210** comparison before requesting data from the victim storage **218**, and thus, the tag RAM **210** comparison between the address of the read operation and the record of cached addresses may extend into a second or third clock cycle.

(610) In the second cycle, the L1 data cache **110** may request the data and ECC syndrome bits from the victim storage **218** if the arbitration manager **1114** permits. In this cycle, the L1 data cache **110** may also determine whether newer data is available in the victim cache store queue **216** by comparing the read address to the pending store address data store **1116**. If so, the priority multiplexer **1118** is set to forward the appropriate data from the victim cache store queue **216**.

(611) Data and ECC may be provided by the victim storage **218** in the third cycle. However, this data may or may not correspond to the memory address specified by the read operation because the

L1 data cache **110** may allocate multiple extended memory addresses to the same entry in the cache's victim storage **218**. Accordingly, in the third cycle, the L1 data cache **110** determines whether the provided data and ECC from the victim storage **218** corresponds to the memory address in the read operation (e.g., a cache hit) based on the comparison of the tag RAM **210** record. In the event of a cache hit, the data and ECC bits are received by the error detection and correction circuit **312**, which corrects any errors in the data in a fourth cycle.

(612) As explained above, newer data that has not yet been written to the victim storage **218** may be present in the victim cache store queue **216**, and may be forwarded from the victim cache store queue **216** by the priority multiplexer **1118**. If so, the MUX circuit **316** selects the forwarded data over the corrected data from the victim storage **218**.

(613) Either the corrected data from the victim storage **218** or the forwarded data from the victim cache store queue **216** is provided to the L1 data cache **110** in a fifth cycle. In this way, the controller **220** may provide data with full ECC checking and correction in the event of a cache hit in about 5 cycles.

(614) In the event that the data and ECC bits are not present in the victim storage **218** (e.g., a cache miss), the victim cache store queue **216** may stall until the data can be retrieved from either the extended memory **106** and/or the victim storage **218**, at which point the data may be written to the victim storage and the tag RAM **210** may be updated.

(615) FIGS. **11B-1** and **11B-2** (collectively FIG. **11B**) illustrate a schematic illustration of an example unified cache store queue **1124**. In the example illustrated in FIG. **11B**, the unified cache store queue **1124** may implement the example main cache store queue **212** and/or the victim cache store queue **216**. In this manner, the topology illustrates example main cache read and write inputs **1128** and example victim cache read and write inputs **1130**. For example, the main cache read and write inputs **1128** may implement the example read and write inputs of the main cache store queue **212** of FIGS. **2** and/or **4**, and the victim cache read and write inputs **1130** may implement the example read and write inputs of the victim cache store queue **216** of FIGS. **2** and/or **11A**. In the example of FIG. **11B**, the main cache read and write inputs **1128** are pipelined separately from the victim cache read and write inputs **1130**. Accordingly, in operation, the main cache read and write inputs **1128** and/or the victim cache read and write inputs **1130** are configured to obtain read and/or write instructions from the CPU **102**. In some examples disclosed herein, the main cache read and write inputs **1128** and the victim cache read and write inputs **1130** may be referred to as inputs of the unified cache store queue **1124**.

(616) The unified cache store queue **1124** includes an example first pipestage (E2) **1132**, an example second pipestage (E3) **1134**, and an example fourth pipestage (E4) **1136**. In this manner, the unified cache store queue **1124** is coupled to the first pipestage (E2) **1132** via example data pipestages **1138** and **1140**, and the main storage **214** and the victim storage **218** of FIG. **2**.

(617) In operation, the CPU **102** transmits a read and/or a write instruction, which enters the unified cache store queue **1124** via the first pipestage (E2) **1132**. In other examples disclosed herein, the unified cache store queue **1124** may obtain a read and/or write instruction from the CPU **102** via the cache controller **220**. At this point, the example L1 cache **110** compares the address of incoming read and/or write instruction with the main cache tag ram **208** and the victim cache tag ram **210**. Additionally, the determination of whether the read and/or write instruction is intended for the main storage **214** or the victim storage **218** is not yet known to the unified cache storage queue **1124**.

(618) In response to a subsequent clock cycle of the CPU **102**, or the cache controller **202**, the read and/or write instruction is transmitted to the third pipestage (E3) **1134**. At this point, the L1 cache **110** determines, or has determined, if the read and/or write instruction is intended for the main storage **214** or the victim storage **218**. Such a determination is transmitted to the unified cache store queue **1124** as a hit and/or miss signal.

(619) In FIG. **11B**, the physical address of the main storage **214** is a function of the CPU **102**

address. The CPU **102** address determines which set of the direct mapped main storage **214** that maps to the CPU **102** address. In FIG. **11B**, the size of the main storage **214** is 32 kilobytes (KB), and the cache line size is 128 bytes, totaling 256 sets. Thus, the physical address of the CPU **102** may range from Address A0 to Address A255.

(620) In FIG. **11B**, the physical address of victim storage **216** is based on the following logic. First, the CPU **102** address is compared with all 16 entries of the victim storage **216**. In the event the CPU **102** address corresponds to a hit in the victim storage **216**, then the location of entry where the CPU **102** transaction hits is the physical address.

(621) However, if the CPU **102** address does not correspond to a hit (e.g., corresponds to a miss), then the replacement policy chooses a location inside the victim storage **216**. Since there are 16 physical address of the victim storage **216**, the CPU **102** address may range from A0 to A15.

(622) In the example of FIG. **11B**, the information corresponding to whether the CPU **102** address is a hit or a miss is sent to the unified cache store queue **1124**. Based on this information, the read and/or write instruction obtained by the CPU **102** enters either the main cache store queue **212** of the unified cache store queue **1124** or the victim cache store queue **216** of the unified cache store queue **1124**.

(623) Example methods, apparatus, systems, and articles of manufacture for read-modify-write support in multi-banked data ram cache for bank arbitration are disclosed herein. Further examples and combinations thereof include the following:

(624) Example 1 includes a data cache system comprising a store queue including a plurality of bank queues, the plurality of the bank queues including a first bank queue, the first bank queue having a write port and a read port, the write port configured to receive a write operation, the read port configured to receive a read operation, storage coupled to the store queue, the storage including a plurality of data banks, the plurality of the data banks including a first data bank, the first data bank having a first port configured to receive the write operation or the read operation, a first multiplexer having a first select input and a first output, the first output coupled to the write port, a second multiplexer having a second select input and a second output, the second output coupled to the read port, a third multiplexer having a third select input and a third output, the third output coupled to the first port, and bank arbitration logic including first arbiters and second arbiters, the first arbiters including a first arbiter and the second arbiters including a second arbiter, the first arbiter coupled to the second arbiter, the second select input, and the third select input, the second arbiter coupled to the first select input.

(625) Example 2 includes the data cache system of example 1, further including first address control logic coupled to a first input of the first arbiter, a first input of the second arbiter, a first input of the first multiplexer, a first input of the second multiplexer, and a first input of the third multiplexer, second address control logic coupled to a second input of the first arbiter, a second input of the second arbiter, a second input of the first multiplexer, a second input of the second multiplexer, and a second input of the third multiplexer, and third address control logic coupled to a third input of the first arbiter, a third input of the second arbiter, a third input of the first multiplexer, a third input of the second multiplexer, and a third input of the third multiplexer.

(626) Example 3 includes the data cache system of example 2, further including a first interface coupled to an input of the first address control logic, the first interface coupled to a scalar data path, a second interface coupled to an input of the second address control logic, the second coupled to a direct memory access data path, and a third interface coupled to an input of the third address control logic, the third interface coupled to a vector data path.

(627) Example 4 includes the data cache system of example 3, further including first write data rotate logic coupled to the first interface and the first input of the first multiplexer, the first input of the second multiplexer, and the first input of the third multiplexer, second write data rotate logic coupled to the second interface and the second input of the first multiplexer, the second input of the second multiplexer, and the second input of the third multiplexer, and third write data rotate logic

coupled to the third interface and the third input of the first multiplexer, the third input of the second multiplexer, and the third input of the third multiplexer.

(628) Example 5 includes the data cache system of example 1, wherein a first output of the first bank queue is coupled to a first input of the first arbiter and a first input of the second arbiter.

(629) Example 6 includes the data cache system of example 1, wherein the plurality of the bank queues is a first quantity and the plurality of the data banks is a second quantity, where the first quantity is the same as the second quantity.

(630) Example 7 includes the data cache system of example 1, wherein the storage is a dynamic random access memory, and the plurality of the data banks has a quantity of 16 data banks, each of the 16 data banks having a data width of 64 bytes.

(631) Example 8 includes an encapsulated data cache system comprising encapsulated memory that includes a plurality of bank queues having a respective read port and a respective write port, and a plurality of data banks having a respective first port, the first port configured to receive a read operation or a write operation, a read port coupled to the encapsulated memory, the read port representative of the respective read port of the bank queues, the read port configured to read first data stored at one or more of the bank queues, a write port coupled to the encapsulated memory, the write port representative of the respective write port of the bank queues, the write port configured to write second data to one or more of the data banks, and a second port coupled to the encapsulated memory, the second port representative of the respective first port of the data banks, the second port configured to read third data from one or more of the data banks or write fourth data to one or more of the data banks.

(632) Example 9 includes the encapsulated data cache system of example 8, further including a multiplexer having a select input and a first output, the first output coupled to the read port, and arbiter logic coupled to the select input, the arbiter logic to assign a first interface of a plurality of interfaces to a first bank queue of the bank queues and a first data bank of the data banks, the first interface to transmit transaction data from the first interface to the read port to invoke the read port to read at least one of the first data from the first bank queue or the third data from the first data bank.

(633) Example 10 includes the encapsulated data cache system of example 9, further including address processing logic coupled to the first interface and the arbiter logic, the address processing logic to identify the first bank queue based on address data included in the transaction data.

(634) Example 11 includes the encapsulated data cache system of example 9, wherein the multiplexer is a first multiplexer, the select input is a first select input, the transaction data is first transaction data, and further including a second multiplexer having a second select input and a second output, the second output coupled to the second port, and the arbiter logic coupled to the second select input, the arbiter logic to assign a second interface of the plurality of the interfaces to a second bank queue of the bank queues and a second data bank of the data banks, the second interface to transmit second transaction data from the second interface to the second port to invoke the second port to read at least one of the first data from the second bank queue or the third data from the second data bank.

(635) Example 12 includes the encapsulated data cache system of example 11, wherein the second port is to read data from the second bank queue and the second data bank during a first time and the read port is to read at least one of the first data from the first bank queue or the third data from the first data bank during the first time.

(636) Example 13 includes the encapsulated data cache system of example 9, wherein the multiplexer is a first multiplexer, the select input is a first select input, the transaction data is first transaction data, the arbiter logic is first arbiter logic, and further including a second multiplexer having a second select input and a second output, the second output coupled to the write port, and second arbiter logic coupled to the second select input, the second arbiter logic to not assign a second interface of the plurality of the interfaces to a second bank queue of the bank queues in

response to the first arbiter logic not assigning the second interface to a second data bank of the data banks.

(637) Example 14 includes the encapsulated data cache system of example 8, further including first arbiter logic coupled to the write port, and second arbiter logic coupled to the read port, the first port, and the first arbiter logic, the second arbiter logic to instruct the first arbiter logic to cause the write port to write the second data to the one or more of the bank queues in response to the bank queues having capacity to store the second data.

(638) Example 15 includes the encapsulated data cache system of example 8, further including first arbiter logic coupled to the write port, and second arbiter logic coupled to the read port, the first port, and the first arbiter logic, the second arbiter logic to instruct the first arbiter logic to cause the write port to write the second data to the one or more of the bank queues in response to at least one of the first data retrieved from the one or more bank queues or the third data retrieved from the one or more data banks.

(639) Example 16 includes a method comprising generating first transaction data based on first instructions from a processor and second transaction data based on second instructions from the processor, the first transaction data indicative of a first quantity of data banks to access for first data, the second transaction data indicative of a second quantity of data banks to access for second data, locating the first data for a first data access operation in a first data bank of at least one of a store queue or storage, locating second data for a second data access operation in a second data bank of at least one of the store queue or the storage, and in response to the first data bank being different from the second data bank, assigning the first transaction data to be transmitted to the first data bank and the second transaction data to be transmitted to the second data bank.

(640) Example 17 includes the method of example 16, further including in response to the first data bank being the same as the second data bank, determining whether the first data access operation or the second data access operation is a read data operation or a write data operation, and in response to the first data access operation being the read data operation and the second data access operation being a write data operation, assigning the first transaction data to be transmitted to the first data bank and not assigning the second transaction data to be transmitted.

(641) Example 18 includes the method of example 16, wherein the first data access operation includes reading a first portion of the first data from one or more of the first data banks and writing a second portion of the first data to one or more of the second data banks, the first data bank included in the one or more of the first data banks or the one or more of the second data banks, and further including in response to locating the first portion in the first data banks of the store queue transmitting the first portion from the store queue to the storage, transmitting a control signal from the store queue to arbiter logic, the arbiter logic coupled to the storage, and in response to the arbiter logic obtaining the control signal, writing the first portion and the second portion to the storage.

(642) Example 19 includes the method of example 16, wherein the first data bank is included in the store queue, and further including, in response to the store queue being full, instructing arbiter logic coupled to the store queue to prevent assigning the first transaction to be transmitted to the first data bank.

(643) Example 20 includes the method of example 16, further including determining a data access operation data size based on the first instructions, determining a data storage address based on the first instructions, mapping the data access operation data size and the data storage address to the first quantity of data banks to be accessed, the first quantity of data banks including a first data bank and a second data bank, and determining that the data access operation includes reading a first portion of the first data from the first data bank and writing a second portion of the first data to the second data bank.

(644) Methods and Apparatus to Facilitate Read-Modify-Write Support in a Victim Cache

(645) In examples disclosed herein, the victim storage **218** includes multiple memory banks, each

bank being 64 bits wide. The victim storage **218** is parallel coupled with the main storage **214**. Additionally, the banks in the victim storage **218** include two 32-bit parity blocks. In examples disclosed herein, 7 bits of ECC syndrome is stored for each of the 32-bit parity blocks. Thus, the overall bank width is 118 bits. In other examples disclosed herein, any suitable bank width may be utilized.

(646) In an example operation, in the event the cache controller **220** transmits a write instruction to the victim storage **218**, and in the event the write instruction is not aligned with a parity block on the victim storage **218**, the cache controller **220** indicates to the victim cache store queue **216** to perform a read-modify-write operation.

(647) In examples disclosed herein, there are two cache systems operating in tandem (e.g., the main storage **214** and the victim cache storage **218**). In examples disclosed herein, the main storage **214** is a direct mapped cache element and the victim cache storage **218** is a fully associative cache storage. Both the direct mapped main storage **214** and the fully associative victim cache storage **218** are protected by an error correcting code (ECC). With this, example applications include reading a line from main storage **214**, the ECC logic would correct the cache line and regenerate the ECC syndrome, and then write the line to victim cache storage **218**. Such an application may utilize two additional clock cycles of the CPU (e.g., one clock cycle for error correction by the ECC logic and another clock cycle for ECC syndrome regeneration). To reduce this latency, examples disclosed herein include utilizing the same parity block size between the main storage **214** and the victim cache storage **218**. Thus, both the main storage **214** and the victim cache storage **218** calculate and/or otherwise determine parity on a 32-bit boundary. In this manner, the L1 cache **110** can move a cache line directly from main storage **214** to the victim cache storage **218** with less latency.

(648) In examples disclosed herein, the L1 data cache **110** supports a number of operations that read data from the cache and make changes to the data before rewriting it. For example, the L1 data cache **110** may support read-modify-write operations. A read-modify-write operation reads existing data and overwrites at least portion of the data. In ECC embodiments, a read-modify-write operation may be performed when writing less than a full bank width. The read functionality of the read-modify-write is used because the portion of the data in the bank that will not be overwritten still contributes to the ECC syndrome bits.

(649) A read-modify-write operation may be split into a write operation and a read operation, and the victim cache store queue **216** may be structured such that the read operation in the read path stays synchronized with the write operation in the victim cache store queue **216**. The read operation and the write operation remain synchronized until the read-modify-write merge component **1108** overwrites at least a portion of the read data with the write data to produce merged data. The merged data is provided to the ECC generator **1112** that generates new ECC syndrome bits for the merged data, and then the merged data and ECC syndrome bits may be provided to the arbitration manager **1114** for storing in the victim storage **218**.

(650) In an example read operation, the L1 data cache **110** compares an address of the read operation to a record of the tag RAM **210**, and the L1 data cache requests the data and ECC syndrome bits from the victim storage **218** and/or the victim cache store queue **216**. Because the read-modify-write operation will modify the data, in examples that track MESI (Modified, Exclusive, Shared, and Invalid) states of entries in the victim storage **218**, a cache hit that is not in either the Modified or Exclusive state may be considered a cache miss. When the data is obtained in the proper state and any errors are corrected, it is provided to the read-modify-write merge component **1108**. In this same cycle, the read-modify-write merge component **1108** may overwrite at least a portion of the corrected data with the write data to produce merged data. The ECC generator **1112** generates new ECC syndrome bits for the merged data.

(651) In an example write operation, the cache controller **220** of FIGS. 2 and/or 3 transmits a write request indicating byte(s) of a word, or an entire word, to be re-written. In such an example, the

write request transmitted by the cache controller **220** includes an address value of the byte and the byte value (e.g., a set of data). For example, the victim storage **218** of FIGS. **2** and/or **3** may include the four-bit word 0x12345678 associated with addresses A3, A2, A1, A0. In such an example, address A3 corresponds to the byte 0x12, address A2 corresponds to the byte 0x34, address A1 corresponds to the byte 0x56, and address A0 corresponds to the byte 0x78 of the stored word. In such an example operation, the cache controller **220** may transmit a write request to replace address A3 with the byte 0x33, replace address A1 with the byte 0x22, and replace address A0 with the byte 0x11 of the currently stored word 12345678. Further in such an example, the first write request to replace address A3 of the stored word with the byte 0x33 would result in the stored word becoming 0x33345678, the second write request to replace address A1 of the stored word with the byte 0x22 would result in the stored word becoming 0x33342278, and the third write request to replace address A0 of the stored word with the byte 0x11 would result in the stored word becoming 0x33342211.

(652) To coordinate such a request, at a first example cycle (e.g., to complete the first write request of replacing address A3 with the byte 0x33), because the ECC generator **1112** operates on word granularity, the cache controller **220** initiates a read request of the currently stored byte (e.g., a read request of a second set of data stored in the victim storage **218**) in address A3 of the currently stored word. In this example, the byte and address in the first write request (e.g., 0x33 and A3) is stored in the latch **1102b**. In addition, the cache controller **220** transmits a read request of the entire currently stored word to the victim storage **218**. Thus, at the end of the first cycle, a read request of the entire currently stored word is transmitted to the victim storage **218** and the byte 0x33 is stored in the first latch **1102b**.

(653) At an example second cycle (e.g., to complete the second write request of replacing address A1 with the byte 0x22) the byte from the first write request is transmitted to the latch **1102c**, the entire currently stored word is transmitted from the victim storage **218** to the ECC logic **312**, and the second write request (e.g., to replace address A1 with the byte 0x22) is transmitted by the cache controller **220** to be stored in the latch **1102b**. During the second cycle, the read-modify-write merge component **1108** obtains the byte stored in the latch **1102c** and the entire currently stored word transmitted by the ECC logic **312**. In this manner, the read-modify-write merge component **1108** identifies the address of the byte in the currently stored word to be updated. Once the read-modify-write merge component **1108** identifies and/or otherwise obtains (a) the value (e.g., byte value, bit value, etc.) of the portion of the currently stored word to be updated from the latch **1102c** and the (b) currently stored word from the ECC logic **312**, the read-modify-write merge component **1108** writes (e.g., replaces, merges, etc.) the portion of the currently stored word (e.g., the second set of data) with the value of the portion of the currently stored word obtained from the latch **1102c** (e.g., the first set of data). For example, the read-modify-write merge component **1108** writes the value of the portion of the word to an address value corresponding to the portion of the word in the word. In some examples disclosed herein, such a merged set of data is provided by the read-modify-write merge component **1108** for writing to the victim storage **218**.

(654) Example methods, apparatus, systems, and articles of manufacture to facilitate read-modify-write support in a victim cache are disclosed herein. Further examples and combinations thereof include the following:

(655) Example 1 includes an apparatus comprising a first storage coupled to a controller, a second storage coupled to the controller and parallel coupled to the first storage, and a storage queue coupled to the first storage, the second storage, and to the controller, the storage queue to obtain a memory operation from the controller indicating an address and a first set of data, obtain a second set of data associated with the address from at least one of the first storage and the second storage, merge the first set of data and the second set of data to produce a third set of data, and provide the third set of data for writing to at least one of the first storage and the second storage.

(656) Example 2 includes the apparatus of example 1, wherein the memory operation is a read

memory operation indicating to read the second set of data associated with the address from at least one of the first storage and the second storage.

(657) Example 3 includes the apparatus of example 2, wherein the storage queue is configured to obtain the read memory operation prior to merging the first set of data and the second set of data.

(658) Example 4 includes the apparatus of example 1, wherein the store queue includes an error correcting code generator configured to generate error correcting code syndrome bits.

(659) Example 5 includes the apparatus of example 1, wherein the first storage is a main storage and the second storage is a victim storage, the main storage and the victim storage located within a level one cache.

(660) Example 6 includes the apparatus of example 1, wherein the controller is a cache controller configured to be coupled to a central processing unit interface to obtain the memory operation from a central processing unit.

(661) Example 7 includes a system comprising a central processing unit configured to generate a memory operation, a first storage coupled to the central processing unit, a second storage coupled to the central processing unit and parallel coupled to the first storage, and a storage queue coupled to the first storage, the second storage, and to the central processing unit, the storage queue to obtain the memory operation from the central processing unit indicating an address and a first set of data, obtain a second set of data associated with the address from at least one of the first storage and the second storage, merge the first set of data and the second set of data to produce a third set of data, and provide the third set of data for writing to at least one of the first storage and the second storage.

(662) Example 8 includes the system of example 7, wherein the memory operation is a read memory operation indicating to read the second set of data associated with the address from at least one of the first storage and the second storage.

(663) Example 9 includes the system of example 8, wherein the storage queue is configured to obtain the read memory operation prior to merging the first set of data and the second set of data.

(664) Example 10 includes the system of example 8, wherein the store queue includes an error correcting code generator configured to generate error correcting code syndrome bits.

(665) Example 11 includes the system of example 10, wherein the error correcting code generator generates the error correcting code syndrome bits in response to merging the first set of data and the second set of data.

(666) Example 12 includes the system of example 8, wherein the first storage is a main storage and the second storage is a victim storage, the main storage and the victim storage located within a level one cache.

(667) Example 13 includes the system of example 8, further including a cache controller configured to be coupled to the store queue, the first storage, the second storage, and the central processing unit via an interface, the cache controller to obtain the memory operation from the central processing unit.

(668) Example 14 includes a method comprising obtaining a memory operation from a controller indicating an address and a first set of data, obtaining a second set of data associated with the address from at least one of a first storage and a second storage, merging the first set of data and the second set of data to produce a third set of data, and providing the third set of data for writing to at least one of the first storage and the second storage.

(669) Example 15 includes the method of example 14, wherein the memory operation is a first memory operation, wherein the address is a first address, further including obtaining a second memory operation indicating a second address and a fourth set of data, and merging the fourth set of data with the third set of data to produce a fifth set of data.

(670) Example 16 includes the method of example 15, further including storing the first address and the second address in a data store.

(671) Example 17 includes the method of example 15, wherein the first storage is a main storage,

and wherein the second storage is a victim storage.

(672) Example 18 includes the method of example 14, further including generating error detection code based on the third set of data, the error detection code to be stored with the third set of data.

(673) Example 19 includes the method of example 14, further including obtaining a read memory operation indicating to read the second set of data.

(674) Example 20 includes the method of example 14, wherein the controller is a cache controller configured to be coupled to a central processing unit interface to obtain the memory operation from a central processing unit.

(675) Methods and Apparatus to Reduce Read-Modify-Write Cycles for Non-Aligned Writes in Victim Cache

(676) When issuing a write instruction to a multi-banked storage, such a write instruction may be transmitted with a corresponding read instruction, regardless of the size of the write instruction, in an attempt to execute a full read-modify-write cycle of such a write instruction. For example, in such an application, a write instruction may be obtained by a CPU indicating to write 128 bits across two 64-bit memory banks, starting at address A0 of the first memory bank. In such an example, though the entirety of both the example memory banks are being written, such an application maintains a read instruction to read the data currently stored in the two example memory banks. However, such an approach is inefficient as twice the processing power (e.g., a write and a read instruction) is needed. Additionally, such an approach does not provide any control logic and/or processing circuitry to analyze the write instruction.

(677) In the example illustrated in FIG. 11A, the main storage **214** and/or the victim storage **218** may be multi-banked storages. For example, the victim storage **218** may include sixteen memory banks (e.g., sixteen sub-RAMs), each 64 bits wide. In such an example, in the event the cache controller **220** transmits a write instruction to write all 64 bits of a first bank of the victim storage **218** (e.g., write a 64-bit word starting with the first address of the first bank), the write instruction can be executed without initiating a read instruction. For example, the bank processing logic **303** may detect that such a write of an entire bank is to be performed and, thus, indicate to the cache controller **220** to initiate the read-modify-write operation, negating to transmit the read instruction.

(678) Similarly, in the event the cache controller **220** transmits a write instruction to write all 128 bits of a first bank and a second bank of the victim storage **218** (e.g., a write instruction indicating to write a 128 bit word starting with the first address of the first bank and ending with the last address of the second bank), the write instruction can be implemented without initiating a read instruction. For example, the bank processing logic **303** may detect that such a write of the entirety of multiple banks is to be performed and, thus, indicate to the cache controller **220** to initiate the read-modify-write operation, negating to transmit the read instruction.

(679) However, in some examples disclosed herein, the cache controller **220** may transmit a write instruction to write 130 bits of a first bank, a second bank, and a third bank of the victim storage (e.g., a write instruction indicating to write a 130 bit work starting with the first address of the first bank and ending with the second address of the third bank). In such an example, the bank processing logic **303** detects that all addresses of the first bank and the second bank of the victim storage **218** are to be written entirely and, thus, indicate to the cache controller to initiate the read-modify-write operations for the first bank and the second bank of the victim storage, negating to transmit the read instruction. In such an example, the bank processing logic **303** may detect that the third bank of the victim storage **218** is to be partially written (e.g., two addresses of the 64 addresses are to be written), and, thus, indicate to the cache controller **220** to initiate a full read-modify-write operation of the third bank of the victim storage **218**. Example description of a read-modify-write operation is described above.

(680) Aggressive Write Merging to Reduce Bank Pressure in Victim Cache

(681) In the example of FIG. 11A, the example victim cache store queue **216** stores a number of write operations at different sections of the victim cache store queue **216** (e.g., at the example

latches **1102a-e**). For example, when the CPU **102** transmits three separate write operations in a row, the first write operation that the CPU **102** provided is stored at the first latch **1102b** and moved to the second latch **1102c** when the second operation is received at the first latch **1102b**. Accordingly, after the three write operations are received, the first latch **1102b** will store and/or output the last write operation with respect to time (e.g., which is last to be stored in the victim storage **218**), the second latch **1102c** will have the second write operation (e.g., which is second to be stored in the main storage **214**), and the third latch **1102d** will have the first write operation (e.g., which was the first to be stored in the example victim storage **218**). Whenever data in the victim cache store queue **216** is next to be stored in the victim storage **218**, the example arbitration manager **1114** reserves a cycle for the data to be written into the example victim storage **218**. Accordingly, during the reserved cycle, the victim storage **218** may not be available to cannot perform read operations. However, if the data operations stored in two or more of the latches **1102b**, **1102c**, **1102d** correspond to the same memory address, the data can be merged in order to write the data into the memory address of the victim storage **218** once, instead of two or three times. For example, if the write operation stored in the latch **1102d** corresponds to writing a byte of the memory address and the write operation stored in the latch **1102c** corresponds to writing a different byte to the memory address, the second write will overwrite the first write.

(682) Instead of reserving a cycle to write the first byte into the victim storage **218**, which will be overwritten in a subsequent cycle, the victim cache store queue **216** merges the two writes into one write, so that only one cycle is used to write the second transaction (e.g., to avoid reserving a cycle for the first write). Such an aggressive merge reduces the number of cycles reserved for write operations. In this manner, the victim storage **218** will have extra cycles to perform read operations, thereby decreasing the latency of the overall systems.

(683) To enable an aggressive merge of write operations to the same memory address. the output of the example latches **1102b-1102d** are coupled to the example merging circuits **1103a-403c**. For example, the output of the third latch **1102d** may be coupled to the merging circuit **1103a**, the output of the second latch **1102c** may be coupled to the merging circuit **1103b**, and the output of the first latch **1102b** may couple to the merging circuit **1103c**. The output of the merging circuit **1103a** may additionally be coupled to the output of the second latch **1102c** and the merging circuit **1103b**, the merging circuit **1103b** may be coupled to the merging circuit **1103c**, and the merging circuit **1103c** may be coupled to the input of the first latch **1102b**. Additionally or alternatively, there may be additional comparator circuits that compare outputs of a different set of latches (e.g., a comparator circuit to compare the output of latch **1102b** with the output of latch **1102d**). The example merging circuits **1103a-c** include example comparator(s) **1120** and example switches **1122**.

(684) The comparator(s) **1120** compare the memory address locations for each write operation that is stored in the respective latches **1102b-1102d** to determine whether any of the write operations in the example store queue correspond to the same memory address. The example comparator **1120** may be one comparator to compare all the write operations of the latches **1102b-1102d** or may be separate comparators **1120**, to compare two of the latches **1102b-d** (e.g., a first comparator to the memory address of latch **1102b** to the memory address of latch **1102c**, a second comparator to the memory address of **1102b** to the memory address of latch **1102d**, etc.). The comparator(s) **1120** output the results of the comparisons (e.g., with one or more signals corresponding to the one or more comparisons) to the example switch(es) **1122** and/or the arbitration manager **1114**. If the example arbitration manager **1114** receives a signal indicative of a match, the arbitration manager **1114** will not reserve the cycle for a first write operation while the first write operation is merged with a second write operation to the same memory location (e.g., to free up cycles for other cache operations).

(685) The example switch(es) **1122** reroute the write operations in the example latches **1102b-1102d** based on the comparison. For example, if the memory address of the write operation stored

in the example latch **1102d** is the same as the memory address stored in the latch **1102c**, the example switch(es) **1122** enable and/or disable to reroute the output of the latch **1102d** to latch **1102c**, instead of routing to the example arbitration manager **1114**. In this manner, the two write operations are combined and written into the victim storage **218** in a subsequent cycle as a single write operation instead of two write operations. The switch(es) **1122** may be electrical switches, transistors (e.g., MOSFETS), demultiplexers, and/or any other component that can reroute a signal in a circuit.

(686) When a latch receives a rerouted write operation, a MUX of the one of the merging circuits **403a-c** (e.g., as further described above in conjunction with FIG. 4D) performs a merging protocol for the one or more rerouted write operations that prioritizes the newest write operation. For example, if the comparator(s) **1120** determines that the write operation stored in the example latch **1102c** corresponds to the same memory address as the write operation stored in the example latch **1102d**, the switches(es) **1122** reroute the write operation stored in the example latch **1102d** to the latch **1102c**. In such an example, if the write operation in the example latch **1102c** corresponds to writing data to byte0 and byte2 of the memory address and the write operation from the latch **1104d** corresponds to writing data to byte0 and byte1, the example merging circuit **1103a** merges the two write operations to keep the writing data stored in latch **1102c** (e.g., the write to byte0 and byte2) and include the write data from latch **1102d** that doesn't overlap (e.g., byte2). In this example, the write data of byte 0 from the latch **1104d** is discarded because the data to be written at byte 0 from the latch **1104d** will be overwritten by the write instructions of the latch **1102c**. Accordingly, in the above example, the merged data corresponds to the write data for byte0 from latch **1102c**, the write data for byte1 from latch **1104d**, and the write data for byte2 from the latch **1102c**. When a subsequent write operation is added to the victim cache store queue **216**, the merged write data from the latch **1102c** may be manipulated (e.g., via one of the example blocks **1104**, **1106**, **1108**) and/or pushed to the next latch **1102d** to be stored in the example victim storage **218** during a subsequent cycle. An example hardware implementation of the merging protocol is further described above in conjunction with FIG. 4C.

(687) Optimized Atomic Operations and Histogram Operations Support in L1 in Victim Cache

(688) Atomic operations are further example of multi-part memory operations. For example, an atomic compare and swap operation manipulates a value stored in the memory location based on the results of a comparison of the existing value stored at the memory location. For example, the CPU **102** may want to replace the data stored in the L1 cache **110** with a new value if the existing value stored in the L1 cache **110** matches a specific value.

(689) In some example systems, when a CPU wanted to perform an atomic operation, the CPU sent a read operation to a memory address, performed the manipulation on the read data, and then executed a write operation to the same memory address to store the manipulated data. Additionally, in such systems, the L1 cache may need to paused, rejected, blocked, and/or halted any transactions from other devices (e.g., other cores of the CPU, higher level cache, the extended memory, etc.) until the atomic operation was complete (e.g., to avoid manipulation of the memory address corresponding to the atomic operation during the atomic operation). Accordingly, such example techniques may require lots of effort on behalf of the CPU and lots of reserved cycles that increase latency.

(690) In the example of FIG. 11A, the example victim cache store queue **216** handles atomic operations in conjunction with the read modify write structure. In this manner, the example CPU **102** can send a single atomic operations operation to the L1 cache **110**, and the victim cache store queue **216** handles the atomic data manipulation and writing operation. Accordingly, the CPU **102** utilizes a single cycle to execute an atomic operation and can use the other cycles (e.g., used in atomic protocols) to perform other functions, thereby reducing the latency of the overall computing system **100**.

(691) In operation, when the CPU **102** transmits an atomic operation and/or an atomic compare and

swap operation to increment and/or swap the data at a memory address by a value of 1, for example, the atomic instruction is received by the latch **1102a** and the tag RAM **210** verifies whether the memory address is stored in the example victim storage **218**. If the memory address is stored in the example victim storage **218**, the tag RAM **210** instructs the example victim storage **218** to output the data at the memory address while the atomic instructions are passed to the example latch **1102b**. While the victim storage **218** outputs the data to the latch **324a**, the example latch **1102b** outputs the atomic operation to the latch **1102c**. At this point the ECC logic **312** performs error detection and/or correction protocol as described above, and the data from the memory address location is forwarded to the example arithmetic component **1104** (e.g., for atomic operations) or the atomic compare component **1106** (e.g., for the atomic compare and swap operations). The arithmetic component **1104** obtains the atomic operation (e.g., including data identifying how to manipulate the data) and/or the atomic compare and swap **1106** obtains the atomic compare and swap operation (e.g., including a key and data to be written if the key matches read data) from the latch **1102c** and obtains the data from the corresponding memory address from the output of the ECC logic **312**. The arithmetic component **1104** performs the manipulation to the data (e.g., increment the data by 1) and/or the atomic compare component **1106** may perform the swap (replaces the data if the read data matches a key, etc.) and outputs the incremented and/or swapped-in and outputs the incremented data for the corresponding memory address (e.g., the atomic result) to the example latch **1102d** via the example MUX **1110**. (e.g., which is enabled via the cache controller **220**). The latch **1102d** outputs the new data corresponding to the memory address to the ECC generator **1112** to generate the ECC bit and the arbitration manager **1114** writes the new data (e.g., the atomic result and/or atomic compare and swap result) to the memory address in conjunction with the ECC bit in the example victim storage **218**. Additionally or alternatively, the corrected value out of the EDD logic **1112** is returned to the CPU **102**. Thus, the atomic operation is performed with only one instruction from the CPU **102**.

(692) The atomic compare component **1106** and/or the arithmetic component **1104** have several inputs. For example, the atomic component **1106** receives (e.g., obtains) the type of atomic to perform (e.g. atomic compare and swap, or atomic swap), the new data to swap in, the ECC corrected data read out the of the cache **310**, and the size of the size of the data to be manipulated during the atomic operation (e.g., 32-bit or 64-bit), In the example circuit implementation **450** of FIG. **4B**, the atomic compare component **1106** receives an atomic compare and swap operation and the arithmetic component **1104** receives an atomic operation. The atomic compare component **1106** compares the comparison value (e.g., a key) provided by the CPU **102** against the ECC data **310**. On a match, the new data is swapped in place of the old data (e.g. ECC data **310**) and output to the MUX **1110**. The size of the new data swapped-in is determined by `cas_acc_sz` input (e.g. 32-bit or 64-bit). In the example circuit implementation **450** of FIG. **4C**, the atomic compare component **1106** may also receive an atomic swap operation. The atomic compare component **1106** will swap-in the new data replacing the ECC data **310** regardless of the comparison result and output the new value to the mux **1110** and the old data from the address is read from the main storage **214** and is provided back to the CPU **102**. The size of the new data swapped-in is determined by `cas_acc_sz` input (e.g. 32-bit or 64-bit). In the example circuit implementation **450** of FIG. **4C**, the arithmetic component **1104** may also receive an atomic operation. The arithmetic component **1104** will manipulate the ECC data **310** and store the manipulated data in the main storage element **214**. The size of the new data swapped-in is determined by `cas_acc_sz` input (e.g. 32-bit or 64-bit). A histogram operation is where the CPU **102** wants to know the value of a bin stored many of each value is present in a section of victim storage **218** (e.g., a SRAM line from the SRAM portion of the victim storage **218**). For example, if a SRAM line has 6 bins with the first bin storing 0, the second bin storing 0, the third bin storing 2, the fourth bin storing 0, the fifth bin storing 0, and the sixth bin storing 3. A histogram of the SRAM line may correspond to [0, 0, 2, 0, 0, 3].

Alternatively, the histogram may be structured in a different manner (e.g., [3, 0, 0, 2, 0, 0]). In some

example systems, to perform a histogram function, the CPU has to read each individual value and increment for each value.

(693) For example, to determine how many 0s are stored in a 10 byte SRAM line, the CPU will perform 10 reads. Then to determine how many 1s are in the same 10 byte SRAM line with 10 bins, the CPU will perform an additional 10 reads. Thus, to perform a histogram operation such system may need to perform $(N)(M)$ reads, where N is the size of the section of memory (e.g., 10 bytes) being read and M is the number of values that could be stored in each byte. Like an atomic operation, the L1 SRAM may have to block, pause, halt, discard, etc. all other read and/or write operations until the histogram operation is complete.

(694) In the example of FIG. 11A, the CPU **102** instructs the victim storage **218** to perform the histogram operation. Thereby changing the number of cycles that the CPU **102** has to reserve for the operation from $(N)(M)$ to 1. Additionally, because the atomic operation protocol is already implemented in the store queue, the histogram operation can be performed using the arithmetic component **1104** by performing N reads for the N size of the memory and incrementing a count for each value in the example victim SRAM store queue **216**, thereby reducing the number of read operation from $(N)(M)$ operations to N operations.

(695) In operation, when the CPU **102** transmits a histogram operation corresponding to a section (e.g., a SRAM line) of the victim storage **218**, the operation is stored in the example latch **1102a** while the tag RAM **210** verifies whether the memory address corresponding to the histogram operation is available in the victim storage **218**. The example cache controller **220** facilitates the read operation for each byte of the section identified in the histogram operation (e.g., where histogram bins are accessed in parallel by reading up to 128 Bytes at the same time). If available, the tag RAM **210** instructs the victim storage **218** to output the data at a first byte of the section of the victim storage **218** while the histogram operation is output by the example latch **1102a** to the example latch **1102b**. When the example victim storage **218** outputs the data that has been read from the memory address to the example latch **324a**, the latch **1102b** outputs the histogram operation to the example latch **1102c**. After the ECC logic **312** performs the error detection and correction functionality, the data read at the byte is sent to the example arithmetic component **1104**.

(696) After receiving the read value from the ECC logic **312** and the histogram instructions from the latch **1102c**, the arithmetic component **1104** initiates data representative of the histogram. For example, the arithmetic component **1104** may initiate a vector (e.g., representing a histogram) with an initial value (e.g., zero) for each possible value that could be stored in the bytes of the victim storage. The arithmetic component **1104** increments the value of the vector based on output by the ECC logic **312** (e.g., the read byte). For example, if the read value of the byte is 0, the arithmetic component **1104** increments the value corresponding to 0 in the vector.

(697) Because each vector is incremented in parallel, the resulting vector corresponds to a histogram of the values that were read in the corresponding sections of SRAM in parallel. Because a value of the histogram is incremented for each bit, the resulting vector is a histogram of the values stored in the section of memory identified in the histogram operation from the CPU **102**. In some examples, the arithmetic component **1104** may increment in parallel by some weighted value (e.g., 1.5). When complete, the example histogram is input to the example MUX **418** (e.g., controlled by the example pending store address table **1116**) to be input to the MUX **316** via the example latch **1102e**. Once the operation is complete, the example cache controller **220** controls the MUX **316** to output the final histogram vector to the example CPU interface **202** via the multiplexer circuit **314** and the example latch **322b**, thereby ending the histogram operation.

(698) In some examples, the L1 cache **110** supports functionality where a histogram bin can saturate after the histogram bin includes more than a threshold limit of the bin size (e.g., a byte, a halfword, a word, etc.). The above Table 1 illustrates an example of saturation values. Using this functionality, the histogram bin values will not roll over once they reach the maximum value.

(699) Example methods, apparatus, systems, and articles of manufacture to facilitate victim cache

and atomic operations are disclosed herein. Further examples and combinations thereof include the following: Example 1 includes a system comprising a first cache storage to store a first set of data, a second cache storage to store a second set of data that has been evicted from the first cache storage, and a storage queue coupled to the first cache storage and the second cache storage, the storage queue including an arithmetic component to receive the second set of data from the second cache storage in response to a memory operation, and perform an arithmetic operation on the second set of data to produce a third set of data, and an arbitration manager to store the third set of data in the second cache storage.

(700) Example 2 includes the system of example 1, wherein the first cache storage is a main storage and the second cache storage is a victim storage.

(701) Example 3 includes the system of example 1, wherein the arithmetic component to obtain (a) the second set of data from the second cache storage via an error detection and correction circuit and (b) the memory operation from a central processing unit via a latch.

(702) Example 4 includes the system of example 1, wherein the third set of data is stored in the second cache storage with a single instruction from a central processing unit at a single cycle.

(703) Example 5 includes the system of example 1, further including a modified, exclusive, shared, invalid (MESI) component to determine a state of a memory address included in the memory operation, and an interface to, if the memory address included in the memory operation corresponds to a shared state, send miss instructions to another cache, the miss instructions including an exclusive state request.

(704) Example 6 includes the system of example 5, wherein the arbitration manager is to store the third set of data at the memory address in the second cache storage after the exclusive state request has been granted from the other cache.

(705) Example 7 includes the system of example 5, wherein the arithmetic component is to receive the second set of data from the second cache storage after the exclusive state request has been granted from the other cache.

(706) Example 8 includes the system of example 5, wherein the second cache storage and the first cache storage are connected in parallel to a central processing unit.

(707) Example 9 includes the system of example 5, wherein the memory operation is an atomic operation.

(708) Example 10 includes a storage queue comprising an arithmetic component to receive a second set of data from a cache storage in response to a memory operation, and perform an arithmetic operation on the second set of data to produce a third set of data, and an arbitration manager to store the third set of data in the cache storage.

(709) Example 11 includes the storage queue of example 10, wherein the cache storage is a victim cache storage, the victim cache storage storing data that has been removed from a main cache storage.

(710) Example 12 includes the storage queue of example 11, wherein the victim cache storage and the main cache storage are connected in parallel to a central processing unit.

(711) Example 13 includes the storage queue of example 10, wherein the arithmetic component to obtain (a) the second set of data from the cache storage via an error detection and correction circuit and (b) the memory operation from a central processing unit via a latch.

(712) Example 14 includes the storage queue of example 10, wherein the third set of data is stored in the cache storage with a single instruction from a central processing unit at a single cycle.

(713) Example 15 includes the storage queue of example 10, wherein the arbitration manager is to store the third set of data at a memory address in the cache storage after an exclusive state request has been granted from another other cache.

(714) Example 16 includes the storage queue of example 15, wherein the arithmetic component is to receive the second set of data from the cache storage from the memory address after the exclusive state request has been granted from the other cache.

(715) Example 17 includes the storage queue of example 10, wherein the memory operation is an atomic operation.

(716) Example 18 includes a method comprising obtaining a second set of data from a cache storage in response to a memory operation, and performing an arithmetic operation on the second set of data to produce a third set of data, and storing store the third set of data in the cache storage.

(717) Example 19 includes the method of example 18, wherein the cache storage is a victim cache storage, the victim cache storage storing data that has been removed from a main storage.

(718) Example 20 includes the method of example 19, further including storing the third set of data at a memory address in the cache storage after an exclusive state request has been granted from another other cache

(719) Atomic Compare and Swap Support in L1 in Victim Cache for Coherent System

(720) If data from an address is stored in the L1 cache **110**, data for the same address may also be cached in other caches (e.g., L2 cache **112**, and/or L3 cache **114**). Accordingly, the memory address needs to be tracked and/or monitored to ensure that any changes to the data in one cache is identified, tracked, and/or reflected to the other caches (e.g., L2 cache **112** and L3 cache **114**). As described above in conjunction with FIGS. 3A-3D, the example MESI RAM **300** tracks the state of the data stored in the victim storage **218** to be able to avoid issues with mismatched data in different caches that correspond to the same memory address. For example, if the CPU **102** transmits a read operation, the example MESI RAM **300** changes the state of the memory address to shared, because the data in the memory address will not be manipulated. If the CPU **102** transmits a write operation, the example MESI RAM **300** changes the state of the memory address to exclusive, because the data in the memory address will be manipulated and the victim storage **218** needs write permission for the address. After the data in the memory address is written to the victim storage **218**, the MESI RAM **300** updates the state of the memory address to modified (e.g., indicating that the memory address has been modified).

(721) As described above, during atomic and/or histogram protocols, the data from a memory address is read from the victim storage **218** and provided to the victim cache store queue **216** to be updated (e.g., incremented) and written back into the victim storage **218**. However, if the MESI RAM **300** has identified the state of the corresponding memory address as in shared state, the write operation of the atomic protocol may cause problems with other level caches (e.g., because the write will cause a mismatch of data in different caches).

(722) To avoid such data mismatches during atomic operations (e.g., including atomic compare and swap), the example cache controller **220** marks cache hits that correspond to a shared state as a cache miss. In this manner, the cache controller **220** can instruct the L2 interface **228** to send the cache miss to the higher level cache with an exclusive state request. In this manner, the higher level cache can grant the exclusive state to the L1 cache **110** and the L1 cache **110** can perform the read and write operation as part of the atomic operation in response to receiving the granted exclusive state.

(723) After the cache miss with exclusive request is sent to the higher level cache and the higher level cache responds with the write data and the exclusive response, the example atomic operation logic **1106** will instruct the MESI RAM **300** to tag the data as modified. The received data from the L2 cache **112** is transmitted into the victim cache store queue **216** to be stored in the victim storage **218**. Because the operation was an atomic operation (e.g., a regular atomic operation or an atomic compare and swap) or a histogram protocol, the data from the higher level cache is manipulated by the example arithmetic component **1104** and/or the example atomic compare component **1106** for the manipulation and stored in the example victim storage **218** via the example ECC generator **1112** and the example arbitration manager **1114**.

(724) While an example manner of implementing the L1 data cache **110** of FIG. 1 is illustrated in FIGS. 2-5 and/or 10-11, one or more of the elements, processes and/or devices illustrated in FIGS. 2-5 and/or 10-11 may be combined, divided, re-arranged, omitted, eliminated and/or implemented

in any other way. Further, the example CPU interface **202**, the example tag RAM accesses **204**, **206**, the example tag RAM **208**, **210**, the example main cache store que **212**, the example main storage **214**, the example victim cache store queue **216**, the example victim storage **218**, the example cache controller **220**, the example main cache controller **222**, the example victim cache controller **224**, the example L2 interface **228**, the example MESI RAM **300**, the example address processing components **302a-c**, the example bank processing logic **303**, the example hit-miss comparison logic **304**, **306**, the example replacement policy component **208**, the example ECC logic **310**, the example MUX circuit **314**, **316**, **318**, **320**, **330**, **330a-c**, **410**, **418**, **508** the example latches **322a**, **324a**, **402a-d**, **1102-a-d**, the example address encoder **326**, **326a**, **326b**, **328**, **328a**, **328b** the example address read **332**, the example merging circuits **403a-c**, **1103a-c**, the example arithmetic component **404**, **1104** the example atomic compare component **406**, **1106**, the example RMW component **408**, **1108**, the example ECC generator **412**, **1112**, the example arbitration manager **414**, **1114**, the example pending store address database **416**, **1116**, the example comparator **420**, **1120**, the example switches **422**, **1122**, the example ports **424**, **426**, **1124**, **1126** the interfaces **502**, **504**, the example address **506**, and/or more generally the L1 data cache **110** of FIGS. 2-5 and/or 10-11 may be implemented by hardware, software, firmware and/or any combination of hardware, software and/or firmware. Thus, for example, any of the example CPU interface **202**, the example tag RAM accesses **204**, **206**, the example tag RAM **208**, **210**, the example main cache store que **212**, the example main storage **214**, the example victim cache store queue **216**, the example victim storage **218**, the example cache controller **220**, the example main cache controller **222**, the example victim cache controller **224**, the example L2 interface **228**, the example MESI RAM **300**, the example address processing components **302a-c**, the example bank processing logic **303**, the example hit-miss comparison logic **304**, **306**, the example replacement policy component **208**, the example ECC logic **310**, the example MUX circuit **314**, **316**, **318**, **320**, **330**, **330a-c**, **410**, **418**, **508** the example latches **322a**, **322b**, **324a**, **322b**, **402a-d**, **1102-a-d**, the example address encoder **326**, **326a**, **326b**, **328**, **328a**, **328b** the example address read **332**, the example merging circuits **403a-c**, **1103a-c**, the example arithmetic component **404**, **1104** the example atomic compare component **406**, **1106**, the example RMW component **408**, **1108**, the example ECC generator **412**, **1112**, the example arbitration manager **414**, **1114**, the example pending store address database **416**, **1116**, the example comparator **420**, **1120**, the example switches **422**, **1122**, the example ports **424**, **426**, **1124**, **1126** the interfaces **502**, **504**, the example address **506**, and/or more generally the L1 data cache **110** of FIGS. 2-5 and/or 10-11 could be implemented by one or more analog or digital circuit(s), logic circuits, programmable processor(s), programmable controller(s), graphics processing unit(s) (GPU(s)), digital signal processor(s) (DSP(s)), application specific integrated circuit(s) (ASIC(s)), programmable logic device(s) (PLD(s)) and/or field programmable logic device(s) (FPLD(s)). When reading any of the apparatus or system claims of this patent to cover a purely software and/or firmware implementation, at least one of the example CPU interface **202**, the example tag RAM accesses **204**, **206**, the example tag RAM **208**, **210**, the example main cache store que **212**, the example main storage **214**, the example victim cache store queue **216**, the example victim storage **218**, the example cache controller **220**, the example main cache controller **222**, the example victim cache controller **224**, the example L2 interface **228**, the example MESI RAM **300**, the example address processing components **302a-c**, the example bank processing logic **303**, the example hit-miss comparison logic **304**, **306**, the example replacement policy component **208**, the example ECC logic **310**, the example MUX circuit **314**, **316**, **318**, **320**, **330**, **330a-c**, **410**, **418**, **508** the example latches **322a**, **322b**, **324a**, **324b**, **402a-d**, **1102-a-d**, the example address encoder **326**, **326a**, **326b**, **328**, **328a**, **328b**, the example address read **332**, the example merging circuits **403a-c**, **1103a-c**, the example arithmetic component **404**, **1104** the example atomic compare component **406**, **1106**, the example RMW component **408**, **1108**, the example ECC generator **412**, **1112**, the example arbitration manager **414**, **1114**, the example pending store address database **416**, **1116**, the example comparator **420**, **1120**, the example switches **422**, **1122**, the

example ports **424**, **426**, **1124**, **1126** the interfaces **502**, **504**, the example address **506**, and/or more generally the L1 data cache **110** of FIGS. 2-5 and/or **10-11** is/are hereby expressly defined to include a non-transitory computer readable storage device or storage disk such as a memory, a digital versatile disk (DVD), a compact disk (CD), a Blu-ray disk, etc. including the software and/or firmware. Further still, the example L1 data cache **110** of FIG. 1 may include one or more elements, processes and/or devices in addition to, or instead of, those illustrated in FIGS. 2-5 and/or **10-11**, and/or may include more than one of any or all of the illustrated elements, processes and devices. As used herein, the phrase “in communication,” including variations thereof, encompasses direct communication and/or indirect communication through one or more intermediary components, and does not require direct physical (e.g., wired) communication and/or constant communication, but rather additionally includes selective communication at periodic intervals, scheduled intervals, aperiodic intervals, and/or one-time events.

(725) A flowchart representative of example hardware logic, machine readable instructions, hardware implemented state machines, and/or any combination thereof for implementing the L1 data cache of FIGS. 1-5 and/or **10-11** is shown in FIGS. 12-33. The machine readable instructions may be one or more executable programs or portion(s) of an executable program for execution by a computer processor such as the processor **3412** shown in the example processor platform **3400** discussed below in connection with FIG. 34. The program may be embodied in software stored on a non-transitory computer readable storage medium such as a CD-ROM, a floppy disk, a hard drive, a DVD, a Blu-ray disk, or a memory associated with the processor **3412**, but the entire program and/or parts thereof could alternatively be executed by a device other than the processor **3412** and/or embodied in firmware or dedicated hardware. Further, although the example program is described with reference to the flowchart illustrated in FIG. 34, many other methods of implementing the example L1 cache **110** may alternatively be used. For example, the order of execution of the blocks may be changed, and/or some of the blocks described may be changed, eliminated, or combined. Additionally or alternatively, any or all of the blocks may be implemented by one or more hardware circuits (e.g., discrete and/or integrated analog and/or digital circuitry, an FPGA, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware.

(726) The machine readable instructions described herein may be stored in one or more of a compressed format, an encrypted format, a fragmented format, a compiled format, an executable format, a packaged format, etc. Machine readable instructions as described herein may be stored as data (e.g., portions of instructions, code, representations of code, etc.) that may be utilized to create, manufacture, and/or produce machine executable instructions. For example, the machine readable instructions may be fragmented and stored on one or more storage devices and/or computing devices (e.g., servers). The machine readable instructions may require one or more of installation, modification, adaptation, updating, combining, supplementing, configuring, decryption, decompression, unpacking, distribution, reassignment, compilation, etc. in order to make them directly readable, interpretable, and/or executable by a computing device and/or other machine. For example, the machine readable instructions may be stored in multiple parts, which are individually compressed, encrypted, and stored on separate computing devices, wherein the parts when decrypted, decompressed, and combined form a set of executable instructions that implement a program such as that described herein.

(727) In another example, the machine readable instructions may be stored in a state in which they may be read by a computer, but require addition of a library (e.g., a dynamic link library (DLL)), a software development kit (SDK), an application programming interface (API), etc. in order to execute the instructions on a particular computing device or other device. In another example, the machine readable instructions may need to be configured (e.g., settings stored, data input, network addresses recorded, etc.) before the machine readable instructions and/or the corresponding program(s) can be executed in whole or in part. Thus, the disclosed machine readable instructions

and/or corresponding program(s) are intended to encompass such machine readable instructions and/or program(s) regardless of the particular format or state of the machine readable instructions and/or program(s) when stored or otherwise at rest or in transit.

(728) The machine readable instructions described herein can be represented by any past, present, or future instruction language, scripting language, programming language, etc. For example, the machine readable instructions may be represented using any of the following languages: C, C++, Java, C#, Perl, Python, JavaScript, HyperText Markup Language (HTML), Structured Query Language (SQL), Swift, etc.

(729) As mentioned above, the example processes of FIGS. 12-33 may be implemented using executable instructions (e.g., computer and/or machine readable instructions) stored on a non-transitory computer and/or machine readable medium such as a hard disk drive, a flash memory, a read-only memory, a compact disk, a digital versatile disk, a cache, a random-access memory and/or any other storage device or storage disk in which information is stored for any duration (e.g., for extended time periods, permanently, for brief instances, for temporarily buffering, and/or for caching of the information). As used herein, the term non-transitory computer readable medium is expressly defined to include any type of computer readable storage device and/or storage disk and to exclude propagating signals and to exclude transmission media.

(730) “Including” and “comprising” (and all forms and tenses thereof) are open ended terms. Thus, whenever a claim employs any form of “include” or “comprise” (e.g., comprises, includes, comprising, including, having, etc.) as a preamble or within a claim recitation of any kind, additional elements, terms, etc. may be present without falling outside the scope of the corresponding claim or recitation. As used herein, when the phrase “at least” is used as the transition term in, for example, a preamble of a claim, it is open-ended in the same manner as the term “comprising” and “including” are open ended. The term “and/or” when used, for example, in a form such as A, B, and/or C refers to any combination or subset of A, B, C such as (1) A alone, (2) B alone, (3) C alone, (4) A with B, (5) A with C, (6) B with C, and (7) A with B and with C. As used herein in the context of describing structures, components, items, objects and/or things, the phrase “at least one of A and B” refers to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B. Similarly, as used herein in the context of describing structures, components, items, objects and/or things, the phrase “at least one of A or B” refers to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B. As used herein in the context of describing the performance or execution of processes, instructions, actions, activities and/or steps, the phrase “at least one of A and B” refers to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B. Similarly, as used herein in the context of describing the performance or execution of processes, instructions, actions, activities and/or steps, the phrase “at least one of A or B” refers to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B.

(731) As used herein, singular references (e.g., “a”, “an”, “first”, “second”, etc.) do not exclude a plurality. The term “a” or “an” entity, as used herein, refers to one or more of that entity. The terms “a” (or “an”), “one or more”, and “at least one” can be used interchangeably herein. Furthermore, although individually listed, a plurality of means, elements or method actions may be implemented by, e.g., a single unit or processor. Additionally, although individual features may be included in different examples or claims, these may possibly be combined, and the inclusion in different examples or claims does not imply that a combination of features is not feasible and/or advantageous.

(732) FIG. 12 is an example flowchart representative of example machine readable instructions 1200 that may be executed by the example L1 cache 110 of FIGS. 1-5 to perform write miss caching in the example victim storage 218 as described above. Although the instructions of FIG. 12 are described in conjunction with the L1 cache 110 of FIGS. 1-5, the instructions may be described

in conjunction with any type of storage in any type of cache.

(733) At block **1202**, the example cache controller **220** allocates a section of the victim storage **218** for write miss information (e.g., a write miss cache). As described above, the write miss information corresponds to when the CPU **102** transmits write instructions to the example L1 cache **110** to a memory address that is not stored in the storages **214**, **218** of the L1 cache **110** (e.g., so the write instructions are sent to higher level cache for execution). At block **1204**, the cache controller **220** accesses the output of the example hit/miss logic **304**, **306** to determine if a current write operation from the CPU **102** (e.g., received by the cache controller **220**) resulted in a write miss (e.g., the memory address from the write operation is not stored in the example storages **214**, **218**). In some examples, where the operation from the CPU **102** missed the main storage **214** but hit the victim storage **218**, the hit miss comparison logic **304** may transmit a write miss to the example victim storage **218**. In such an example, the victim storage **218** discards the write miss information because the operation hit the victim storage **218**.

(734) If the cache controller **220**, after interfacing with the one or more of the example tag RAMs **208**, **210**, determines that a current write operation from the CPU **102** did not result in a write miss (block **1204**: NO), control returns to block **1204** until a write operation results in a write miss. If the cache controller **220** determines that a current write operation from the CPU **102** results in a write miss (block **1204**: YES), the example cache controller **220** determines if the write miss information corresponds to the same memory address as any write miss information already stored in the allocated section (block **1206**). If the example cache controller **220** determines that the write miss information corresponds to the same memory address as any write miss information already stored in the allocated section (block **1206**: YES), the cache controller **220** instructs the example victim storage **218** to merge the write miss information with the stored write miss information corresponding to the same memory address (block **1208**). The example victim storage **218** merges the two write miss information by overwriting the older write miss information with the most recent write miss information when the most recent write miss information overlaps (e.g., corresponds to the same bytes as) the older write miss information (e.g., discarding the older write miss information that overlaps the more recent write miss information) and maintaining the older write miss information that does not overlap the more recent write miss information.

(735) If the example cache controller **220** determines that the write miss information does not correspond to the same memory address as any write miss information already stored in the allocated section (block **1206**: NO), the cache controller **220** stores, in the example victim storage **218**, the write miss information in the allocated section (block **1210**).

(736) At block **1212**, the example cache controller **220** determines if more than a threshold amount of write miss information has been stored in the allocated section. The threshold amount may be set to the size of the victim cache (e.g., the threshold is satisfied when the allocated section is full), the size of the L2 interface **228** (e.g., if the L2 interface has a 64 byte bandwidth, then the threshold is set to 64 bytes), and/or any other amount set by a user and/or manufacturer. If the example cache controller **220** determines that more than the threshold (e.g., a first threshold) amount of write miss information has not been stored in the allocated section (block **1212**: NO), control returns to block **1204**.

(737) If the example cache controller **220** determines that more than the threshold amount of write miss information has been stored in the allocated section (block **1212**: YES), the cache controller **220** selects a threshold (e.g., a second threshold) amount of write miss information (e.g., the N oldest write miss information stored in the allocated section where N corresponds to the threshold) from the allocated section of the victim storage **218** (block **1214**). The second threshold may correspond to (e.g., be the same as) the first threshold and/or may correspond to the bandwidth of the L2 interface **228** (e.g., if the bandwidth of the L2 interface **228** is 64 bytes, then no more than 64 bytes of write miss data is selected). Additionally or alternatively, the cache controller **220** may proceed to block **1210** when a threshold amount of time has occurred. At block **1216**, the cache

controller **220** causes the example victim storage **218** to remove the selected write miss information from the allocated section. At block **1218**, the example L2 cache interface **228** transmits the selected write miss information to the higher level cache (e.g., the L2 cache **112**). As described above, sending multiple write miss information to utilize more of the bandwidth of the L2 interface **112** results in a more efficient system.

(738) FIG. **13** is an example flowchart representative of example machine readable instructions **1300** that may be executed by the example L1 cache **110** of FIGS. **1-5** to facilitate a read-modify-write operation, in conjunction with the above Section 2. In FIG. **13**, the main cache store queue **212** obtains a write instruction transmitted by the CPU **102** (e.g., transmitted through the cache controller **220**) indicating byte(s) of a word, or an entire word, to be re-written. (Block **1302**). In some examples disclosed herein, the write port **426** may obtain the write instruction transmitted by the CPU **102**.

(739) At block **1304**, the main cache store queue **212** transmits the value of the portion of the word to be rewritten to the latch **402b**. (Block **1304**). In some examples, the latch **402b** transmits the value of the portion of the word to be rewritten to the latch **402c**.

(740) At block **1306**, the main cache store queue **212** stores the address value associated with the location of the portion of the word to be rewritten in the pending store address data store **416**. (Block **1306**). Additionally, the tag ram **208** transmits a read instruction (e.g., a read request) of the entire currently stored word to the main storage **214**. (Block **1308**). At block **1310**, the main cache store queue **212** determines whether there has been a subsequent clock cycle of the CPU **102**, or the cache controller **220**. (Block **1310**). In some examples disclosed herein, the latch **402c** determines whether there has been a subsequent clock cycle of the CPU **102**, or the cache controller **220**. In response to determining that there has not been a subsequent clock cycle of the CPU **102**, or the cache controller **220**, (e.g., the control of block **1310** returns a result of NO), the process waits.

(741) Alternatively, in response to determining that there has been a subsequent cycle of the CPU **102**, or the cache controller **220**, (e.g., the control of block **1310** returns a result of YES), the read-modify-write merge component **408** obtains the value of the portion of the word (e.g., the byte) stored in the latch **402c**. (Block **1312**). Additionally, the read-modify-write merge component **408** obtains the entire currently stored word transmitted by the ECC logic **310**. (Block **1314**). In this manner, the read-modify-write merge **408** identifies the address of the byte in the currently stored word to be updated. Once the read-modify-write merge component **408** identifies and/or otherwise obtains (a) the value (e.g., byte value, bit value, etc.) of the portion of the currently stored word to be updated from the latch **402c** and the (b) currently stored word from the ECC logic **310**, the read-modify-write merge component **408** writes (e.g., replaces) the portion of the currently stored word with the value of the portion of the currently stored word obtained from the latch **402c**. (Block **1316**).

(742) At block **1318**, the main cache store queue **212** generates error detection code based on the word, the error detection code to be stored with the word. (Block **1318**). In some examples disclosed herein the ECC generator **412** generating error detection code based on the word, the error detection code to be stored with the word.

(743) The control of block **1318** may be performed in response to an additional subsequent clock cycle of the CPU **102**, or the cache controller **220**.

(744) In response, the main cache store queue **212** determines whether an additional write instruction is obtained. (Block **1322**). the event the main cache store queue **212** determines another write instruction is obtained (e.g., the control of block **1322** returns a result of YES), the process returns to block **1302**. Alternatively, in the event the main cache store queue **212** determines another write instruction is not obtained (e.g., the control of block **1322** returns a result of NO), the process **1300** may wait until a threshold timeout period occurs, thus ending the process **1300**.

(745) FIG. **14** is an example flowchart representative of example machine readable instructions **1400** that may be executed by the example L1 cache **110** of FIGS. **1-5** to facilitate a read-modify-

write operation for non-aligned writes, in conjunction with the above-Sections 3 and/or 14. At block **1402**, the bank processing logic **303** of FIGS. 3A-3D analyzes the write instructions obtained from the CPU **102** (e.g., the write instructions obtained via the CPU interface **202**). For example, the bank processing logic **303** may operate as initial processing circuitry to determine the nature of the write instruction.

(746) In this manner, the bank processing logic **303** determines the number of memory banks to be written to when executing the write instruction. (Block **1404**). For example, the bank processing logic **303** determines the address locations of the write instruction and, as such, determines the banks of either the main storage **214** or the victim storage **218** that include the corresponding address locations. In response, the bank processing logic **303** determines whether all addresses of a memory bank (e.g., a memory bank included in either the main storage **214** or the victim storage **218**) are to be rewritten. (Block **1406**).

(747) In the event the bank processing logic **303** determines all addresses of a memory bank (e.g., a memory bank included in either the main storage **214** or the victim storage **218**) are to be rewritten (e.g., the control of block **1406** returns a result of YES), the bank processing logic **303** indicates to the CPU **102**, or the cache controller **220**, to execute the write instruction without reading the currently stored values in the memory bank. (Block **1408**). For example, the bank processing logic **303** may identify that addresses A0 to A70 are to be rewritten and, thus, determine that the first memory bank (e.g., a memory bank having addresses A0 to A63) are to be rewritten. Thus, such a first memory bank can be rewritten without reading the currently stored values.

(748) Alternatively, in the event the bank processing logic **303** determines that all addresses of a memory bank are not being rewritten (e.g., the control of block **1406** returns a result of NO), or in response to the execution of the control in block **1408**, the bank processing logic **303** determines whether there are additional memory banks to analyze. (Block **1410**).

(749) For example, to execute the control of block **1410**, the bank processing logic **303** determines whether all memory banks affected by the write instruction have been analyzed. In following the example above, the bank processing logic **303** determines that the memory bank including addresses A64 to A70 have not been analyzed. Thus, in the event the bank processing logic **303** determines that there is an additional memory bank to analyze (e.g., the control of block **1410** returns a result of YES), the process **1400** returns to block **1406**. Alternatively, in the event the bank processing logic **303** determines that there are no additional memory banks to analyze (e.g., the control of block **1410** returns a result of NO), the bank processing logic **303** determines whether another write instruction is obtained. (Block **1412**).

(750) In examples disclosed herein, in the event the bank processing logic **303** determines there is another write instruction (e.g., the control of block **1412** returns a result of YES), the process **1400** returns to block **1402**. Alternatively, in the event the bank processing logic **303** determines that there is not another write instruction (e.g., the control of block **1412** returns a result of NO), the process **1400** may wait until a threshold timeout period occurs, thus ending the process **1400**.

(751) FIG. **15** is an example flowchart representative of example machine readable instructions **1500** that may be executed by the example L1 cache **110** of FIGS. 1-5 to perform an aggressive write merge in the example main cache store queue **212** and/or the victim cache store queue **216**, in conjunction with the above sections 4 and/or 15. Although the instructions of FIG. **15** are described in conjunction with the L1 cache **110** of FIGS. 1-5, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **15** are described in conjunction with the main cache store queue **212** and the main storage **214**. However, the instruction of FIG. **15** can likewise be used in conjunction with the victim cache store queue **216** and the victim storage **218**.

(752) At block **1502**, the example comparator(s) **420** of the example merging circuits **403a-c** obtains write instructions from corresponding latches **402b-d**. As described above, each of the latches **402b-d** include different write instructions from the CPU **102**. At block **1504**, the example

comparator(s) **420** of the example merging circuits **403a-c** compare the memory addresses for the write instructions from the latches **402b-d**. For example, the comparator **420** of the merging circuit **403a** compares the memory address for the write instructions output by the latch **402d** with the write instructions output by the latch **402c**. At block **1506**, the example comparator(s) **420** determine if any two or more write instructions output by the latches **402b-d** correspond to the same memory address. If the comparator(s) **420** determine that any two or more write instructions output by the latches **402b-d** do not correspond to the same memory address (block **1506**: NO), control continues to block **1520**, as further described below.

(753) If the comparator(s) **420** determine that any two or more write instructions output by the latches **402b-d** corresponds to the same memory address (block **1506**: YES), control continues to block **1508**. For each group of write instructions corresponding to the same memory address (blocks **1508-1514**), the one or more of the example merging circuit **1103a-c** that receive(s) write instructions for the same memory address maintain(s) the write data for the byte(s) for the newest write instructions (e.g., the write instructions that were more recently received from the CPU **102**) that overlap the write data for the same byte(s) from older write instructions (block **1510**). At block **1512**, the one or more of the example merging circuit **1103a** that receive(s) write instructions for the same memory address update(s) the write data for byte(s) from older write instructions that do not overlap with write data from the newest write instructions. For example, if the merging circuit **1103a**, storing a write instruction to write byte0 of a memory address, receives rerouted data from latch **402d**, the rerouted data corresponding to a write instruction to write byte0 and byte 1 of the memory address, then the merging circuit **1103a** maintains the write instruction to write to byte0 (e.g., discarding the write instruction to write byte0 from latch **402d**, because the write instruction is older) and updates the write instruction to write to byte1 corresponding to the instruction from latch **402b** (e.g., because the instruction does not overlap with the newest write instructions). At block **1516**, the example switch(es) **422** reroute the merged write instructions the latch with the newest write instructions that also corresponds to the matching memory address (block **1508**).

(754) At block **1518**, the one or more of the merging circuits **402a-c** that rerouted write instructions to be merged flags the data that was rerouted. For example, the one or more merging circuits **403a-c** may transmit a signal to the example arbitration manager **414** and/or the cache controller **220**. In this manner, the arbitration manager **414** and/or the cache controller **220** can avoid reserving a cycle to write the data that has be rerouted into a prior latch for merging. At block **1520**, the latch **402b** determines if an additional write instruction has been received. If the latch **402d** determines that an additional instruction has not been received (block **1520**: NO), control returns to block **1520** until additional write instructions are received. If the latch **402d** determines that an additional instruction has been received (block **1520**: YES), control returns to block **1504**.

(755) FIG. **16** is an example flowchart representative of example machine readable instructions **1600** that may be executed by the example L1 cache **110** of FIGS. **1-5** to perform an atomic operation, as described above in conjunction with the above Sections 5 and 16. Although the instructions of FIG. **16** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **16** are described in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.). However, the instruction of FIG. **16** can likewise be used in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.).

(756) At block **1602**, the cache controller **220** and/or the example latch **402a** obtains an atomic operation from the CPU **102**. When the latch **402a** obtains the atomic operation, the cache controller **220** and and/or the latch **402a** sends the memory address for the atomic operation to the example tag RAM **208** to determine whether the data corresponding to the atomic operation is stored in the example main storage **214**. At block **1604**, the cache controller **220** interfaces with the example hit/miss logic **304** to determine if the memory address corresponding to the atomic

operations is stored in the main storage **214**.

(757) If the cache controller **220** determines that the memory address corresponding to the atomic operation is not stored in the main storage **214** (block **1604**: NO), cache controller **220** interfaces with the example L2 cache interface **228** to submit the atomic miss information to higher level cache (e.g., the L2 cache **112** of FIG. **1**) (block **1606**). In this manner, the example L2 cache **112** can return the corresponding data from the memory address corresponding to be stored in the L1 cache **110** to execute the atomic operation. The L2 cache **112** may have the data corresponding to the memory address stored locally or may obtain the data from the L3 cache **114** and/or the extended memory **110** (e.g., via the L3 cache **114**). At block **1608**, the example arithmetic component **404** of the main cache store queue **212** obtains the data corresponding to the memory address from the L2 cache **112** via the L2 interface **228**. For example, the data may be stored in the example main storage **214**, read, and input to the example arithmetic component **404**.

(758) If the cache controller **220** determines that the memory address corresponding to the atomic operation is stored in the main storage **214** (block **1604**: YES), cache controller **220** causes the example arithmetic component **404** to obtain the data corresponding to the memory address of the atomic operation from the main storage **214** (block **1610**). At block **1612**, cache controller **220** causes the example arithmetic component **404** to perform the atomic operation from the CPU **102** in conjunction with the data from the storage and/or higher level cache that corresponds to the atomic operation. For example, while blocks **1602-1610** occur, the atomic operation is sent to the main cache store queue **212** via the latch **402a**. The atomic operation includes the specifics of the operation (e.g., increment, decrement, etc.). Accordingly, the arithmetic component **404** obtains the atomic operation and the data corresponding to the memory address of the atomic operation. Thus, the arithmetic component **404** can perform the atomic operation (e.g., increment, decrement, etc.) using the obtained data (e.g., that corresponds to the memory address of the atomic operation). At block **1614**, the example cache controller **220** controls the MUX **410** (e.g., via the select line) to ensure that the output of the arithmetic component **404** is output to the example latch **402d**.

Accordingly, the manipulated data (e.g., incremented data, decremented data, etc.) can be passed to the example ECC generation **412** to generate an ECC code for the manipulated data (block **1616**).

(759) At block **1618**, the example ECC generation **412** outputs an ECC code for the manipulated data, the manipulated data, and the memory address location to the example arbitration manager **414**. At block **1620**, the cache controller **220** causes the example arbitration manager **414** to store the atomic output (e.g., the manipulated data) in the main storage **214** at the memory address of the atomic operation.

(760) FIG. **17** is an example flowchart representative of example machine readable instructions **1700** that may be executed by the example L1 cache **110** of FIGS. **1-5** to perform a histogram operation, in conjunction with the above Sections 5 and/or 16. Although the instructions of FIG. **17** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **17** are described in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.). However, the instruction of FIG. **17** can likewise be used in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.).

(761) At block **1702**, the cache controller **220** and/or the example latch **402b** of the main cache store queue **212** and/or the example tag RAM **208** obtains a histogram operation from the CPU **102**. As described above, the histogram operation includes determining a total number of each value stored in a section of memory (e.g., a SRAM line). At block **1704**, the cache controller **220** interfaces with the example hit/miss logic **304** to determine if the memory address corresponding to the histogram operation is stored in the SRAM of the main storage **214**. If the cache controller **220** determines that the memory address corresponding to the histogram operation is stored in the SRAM of the main storage **214** (block **1704**: YES), control continues to block **1710**.

(762) If the cache controller **220** determines that the memory address corresponding to the histogram operation is not stored in SRAM of the main storage **214** (block **1704**: NO), the cache controller **220** interfaces with the example L2 interface **228** to transmit the read miss information to higher level cache (e.g., the example L2 cache **112**) (block **1706**). At block **1708**, the cache controller **220** utilizes the example arbitration manager **414** to obtain the read data from the higher level cache via the L2 interface **228** and stores the data corresponding to the memory address of the histogram operation in the SRAM of the main storage **214**. At block **1710**, the cache controller **220** and/or causes the example arithmetic component **404** to initiate a histogram vector with values to be representative of counts for values stored in the section of the SRAM of the main storage **214**. (763) At block **1712**, the cache controller **220** causes the example SRAM of the main storage **214** outputs the read value of the bins corresponding to the section of SRAM in parallel. The read values are output to the example arithmetic component **404** via the ECC logic **310**. At block **1714**, the cache controller **220** utilizes the example arithmetic component **404** to increment one of the elements of the histogram value based on the read values of the bins. For example, if a read value is '01,' the arithmetic component **404** increments the element that corresponds to the '01' count. At block **1722**, the histogram vector is provided to the example MUX **314** via the example MUX **418** and the example latch **402e** and the example MUX **314** outputs the histogram vector to the CPU **102** via the latch **322b** and the CPU interface **202** (block **1722**). In some examples, the histogram vector is additionally or alternatively, stored in the example main storage **214** via the ECC generator **412** and the arbitration manager **414**.

(764) FIGS. **18A** and **18B** illustrate an example flowchart representative of example machine readable instructions **1800** that may be executed by the example L1 cache **110** of FIGS. **1-5** to perform an atomic compare and swap operation, in conjunction with the above Sections 6 and/or 17. However, the flowchart may be described in conjunction with any atomic operation or a histogram operation. Although the instructions of FIGS. **18A** and **18B** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIGS. **18A** and **18B** are described in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.). However, the instruction of FIGS. **18A** and **18B** can likewise be used in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.).

(765) At block **1802**, the cache controller **220** and/or the example latch **402b** of the example main cache store queue **212** obtains an atomic compare and swap operation with a key from the example CPU **102**. As described above, the atomic compare and swap compares the data at a memory address to a key and performs a write to the memory address with swap data if the previously stored data at the memory address matches the key. At block **1804**, the cache controller **220** interfaces with the example MESI RAM **300** to determine the state of the memory address corresponding to the atomic compare and swap operation. As described above in conjunction with FIGS. **3A-3D**, the MESI RAM **300** tracks the states of the memory addresses (e.g., shared, modified, inactive, or exclusive). At block **1806**, the cache controller **220** interfaces with the example MESI RAM **300** to determine if the state of the memory address is inactive (e.g., the memory addressed is not stored in the L1 cache **110**) or is shared (e.g., stored in the L1 cache **110** and stored in another higher level cache).

(766) If the example cache controller **220** determines that the state of the memory address corresponding to the atomic compare and swap is inactive or shared (block **1806**: YES), the cache controller **220** causes the example L2 interface **228** to submit an atomic misinformation to higher level cache with an exclusive state request (block **1808**). As described above in conjunction with FIG. **4A**, if the memory address is stored in the L1 cache **110** and is shared, writing to the data can cause problems if a second core reads and/or writes the data at a different level of cache, because doing so will lead to inconsistent data for the same memory address in different cache.

Accordingly, the example L2 interface **228** transmits the exclusive state request to let the higher level cache know that the L1 cache **110** will perform an operation on the cache for more than one cycle so that different writes do not occur in different caches to the same memory address.

(767) At block **1810**, the example cache controller **220** causes the example MESI RAM **300** to change the state of the corresponding memory address to exclusive. The MESI RAM **3400** may change the state after submitting the request to higher level cache or after receiving a response from the higher level cache. If the state of the memory address was inactive, the higher level cache will return the data at the memory address, which may be stored in the example main storage **214** and/or input to the example atomic compare component **406**. At block **1812**, the example cache controller **220** causes the example atomic compare component **406** to obtain the data corresponding to the memory address (e.g., from the example main storage **214** and/or from the higher level cache).

(768) If the example cache controller **220** determines that the state of the memory address does not correspond to the atomic compare and swap is inactive or shared (block **1806**: NO), the example atomic compare component **406** obtains the data corresponding to the memory address from the main storage **214** (block **1814**). At block **1816**, the cache controller **220** causes the example atomic compare component **406** determines if the obtained data matches the key (e.g., from the atomic swap and compare operation). If the example atomic compare component **406** determines that the obtained data does not match the key (block **1816**: NO), the cache controller **220** causes the atomic compare component **406** to discard the swap data to be written from the atomic compare and swap (e.g., the data that was to be stored if the obtained data matched the key) (block **1818**). In some examples, atomic compare component **406** outputs the obtained data to rewrite the obtained data back into the main storage **214**. If the example atomic compare component **406** determines that the obtained data matches the key (block **1816**: YES), the cache controller **220** causes the example atomic compare component **406** outputs the swap data to be written to the memory address (e.g., from the atomic compare and swap operation) to the example MUX **410** (Block **1820**).

(769) At block **1822**, the example cache controller **220** controls the MUX **410** (e.g., via the select line) to ensure that the output of the atomic compare component **406** is output to the example latch **402d**. Accordingly, the swapped data can be passed to the example ECC generation **412** to generate an ECC code for the swapped data (block **1824**).

(770) At block **1826**, the example ECC generation **412** outputs an ECC code for the swapped data, the swapped data, and the memory address location to the example to the example arbitration manager **414**. At block **1828**, the cache controller **220** causes the example arbitration manager **414** to store the atomic output (e.g., the manipulated data) in the main storage **214** at the memory address of the atomic operation.

(771) FIG. **19** is an example flowchart representative of example machine readable instructions **1900** that may be executed by the example L1 cache **110** of FIGS. **1-5** to perform in-flight data forwarding and invalidation of write instructions from the CPU **102**, in conjunction with the above-Section 7. Although the instructions of FIG. **19** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache.

(772) At block **1902**, the cache controller **220** obtains a read-miss operation from the tag RAM accesses **204**, **206**. For example, the main cache controller **222** of the cache controller **220** obtains the read-miss operation when the CPU **102** issues a read operation that does not include a corresponding memory address in the main storage element **214** and/or the victim storage element **214**. The cache controller **220** determines a victim to evict from the main storage element **214** (block **1904**). For example, the cache controller **220** creates a victim. The victim corresponds to a memory address and a memory location (e.g., cache line) of the main storage element **214** to remove data from.

(773) The cache controller **220** issues a read-invalidate operation to the store queue **212** (block **1906**). For example, the cache controller **220** sends an operation to the read port **424** of the store

queue in response to receiving the memory address of the victim. The read port **424** obtains an address corresponding to victim (block **1908**). For example, the cache controller **220** sends the victim address to the store queue **212** when issuing a read-invalidate operation.

(774) The data store **416** compares the address of the victim to addresses stored in the data store **416** (block **1910**). For example, the data store **416** maintains a log of the addresses associated with each value stored in any of the latches **402a**, **402b**, **402c**, **402d**, **402e**, and/or any of the merging circuits **403a**, **403b**, and/or **403g**. Additionally, the data store **416** stores the victim address corresponding to the read-invalidate operation. The data store **416** determines if any of the addresses in the data store **416** match the address of the victim (block **1912**). For example, the data store **416** determines if any of the latches **402a-d** include values and/or data corresponding to the victim address. In some examples, the latches **402a-d** store outstanding write addresses. The outstanding write corresponds to a write operation has not been completed (e.g., the data of the write operation has not been fully written into the main storage element **214**). In some examples, the store queue **212** writes data to a location (e.g., a cache line) in the main storage element **214** that an allocation policy selected to be the victim.

(775) If the data store **416** determines that any of the addresses in the data store **416** match the address of the victim (e.g., block **1912** returns a value YES), the priority multiplexer **418** forwards the data corresponding to the matching addresses to the MUX circuit **314** (block **1914**). For example, the data store **416** sends the matching address to the priority multiplexer **418**. The priority multiplexer **418** selects the data and/or the values stored in the latches **402a-d** that store the victim address. The priority multiplexer **418** sends the selected data to the latch **402e** to be forwarded to the MUX circuit **314**. The MUX circuit **314** sends the data to the victim storage element **214** and/or the L2 cache **112**.

(776) The store queue **212** updates the valid tag in the appropriate latches (e.g., **402a-d**) to invalid (block **1916**). For example, the latches that store data for the victim address change the valid tag to indicate that the data is invalid. In this manner, the store queue **212** will discontinue writing that data into the main storage element **214**. Thus, the example read-invalidate operation **1900** corrects any errors that would occur if an outstanding write attempted to write data into a cache line that no longer held the write operation address.

(777) FIG. **20** is an example flowchart representative of example machine readable instructions **2000** that may be executed by the example L1 cache **110** of FIGS. **1-5** to determine whether an address is stored in a victim cache, such as the victim storage **218** of FIG. **2** as described above. Although the instructions of FIG. **20** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **20** are described in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.). However, the instructions of FIG. **20** can likewise be used in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.).

(778) The machine readable instructions **2000** begin at block **2002**, at which the L1 cache **110** receives read address(es) from interface(s). For example, the L1 cache **110** can receive ADP_ADDR_E2_DP0 from the scalar interface **502** of the CPU **102** of FIG. **1**, SNP_ADDR_E2_DP0 from the snoop interface of FIGS. **3** and/or **5**, and/or ADP_ADDR_E2_DP1 from the vector interface **502** of the CPU **102** as depicted in FIG. **8B**.

(779) At block **2004**, the L1 cache **110** compares read address(es) to sets of a multi-bank victim cache tag (VCT) random access memory (RAM). For example, the first comparators **850** of FIG. **8B** can compare a first read address of ADP_ADDR_E2_DP0 to respective addresses stored in the sets **846** of FIG. **8B**. In other examples, the second comparators **852** of FIG. **8B** can compare a second read address of ADP_ADDR_E2_DP1 to respective addresses stored in the sets **846**.

(780) At block **2006**, the L1 cache **110** determines whether at least one of the read address(es) is mapped to one of the sets. For example, one of the first comparators **850** can assert a logic one in

response to the read address matching the set **846** that the one of the first comparators **850** is associated with. In other examples, the second comparators **852** can generate HIT_DP1 based on the comparisons. In other examples, the one of the first comparators **850** can generate a logic low in response to the read address not matching the set **846** that the one of the first comparators **850** corresponds to.

(781) At block **2008**, the L1 cache **110** identifies cache hit(s). For example, the first comparators **850** can generate HIT_DP0 based on the comparisons. In such examples, the first comparators **850** can identify a cache hit in response to at least one bit in the HIT_DP0 bit vector being a logic one (e.g., having a bit value of 1). In some such examples, the first comparators **850** can identify a cache miss in response to none of the bits in the HIT_DP0 bit vector being a logic one. In other examples, the second comparators **852** can generate HIT_DP1 based on the comparisons. In such examples, the second comparators **852** can identify a cache hit in response to at least one bit in the HIT_DP1 bit vector being a logic one. In some such examples, the second comparators **852** can identify a cache miss in response to none of the bits in the HIT_DP1 bit vector being a logic one.

(782) At block **2010**, the L1 cache **110** executes cache hit-miss conversion logic. For example, the first address encoder logic circuit **854** can invoke at least one of the first AND gate **864A**, the third comparator **870A**, or the fourth comparator **872A** of FIG. **8B** to convert a cache hit to a cache miss or vice versa in response to example operating conditions. In other examples, the second address encoder logic circuit **856** can invoke at least one of the second AND gate **864B**, the fifth comparator **870B**, or the sixth comparator **872B** of FIG. **8B** to convert a cache hit to a cache miss or vice versa in response to example operating conditions. An example process that may be used to implement block **2010** is described below in connection with FIG. **21**.

(783) At block **2012**, the L1 cache **110** outputs cache hit address(es) based on the cache hit-miss conversion logic. For example, the first address encoder logic circuit **854** can output HIT_ADDR0 in response to executing cache hit-miss conversion logic. In other examples, the second address encoder logic circuit **856** can output HIT_ADDR1 in response to executing cache hit-miss conversion logic.

(784) At block **2014**, the L1 cache **110** determines whether there additional read address(es) have been received. If, at block **2014**, the L1 cache **110** determines additional read address(es) have been received, control returns to block **2002** to receive the read address(es) from the interface(s). If, at block **2014**, the L1 cache **110** determines no additional read address(es) have been received, the example machine readable instructions **2000** of FIG. **20** conclude.

(785) FIG. **21** is an example flowchart representative of example machine readable instructions **2100** that may be executed by the example L1 cache **110** of FIGS. **1-5** to execute cache hit-miss conversion logic as described above. The example machine readable instructions **2100** of FIG. **21** can be executed to implement block **2010** of FIG. **20**. Although the instructions of FIG. **21** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **21** are described in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.). However, the instructions of FIG. **21** can likewise be used in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.).

(786) The machine readable instructions **2100** of FIG. **21** begin at block **2102**, at which the L1 cache **110** determines whether a new address from a first interface has been written to victim cache in a later pipeline stage. For example, the first decoder **860A** can receive VTA_WR_SET0 at the E2 pipeline stage, which can be representative of the scalar interface **502** of FIG. **5** writing an address to the victim storage **218** at the E3 pipeline stage.

(787) If, at block **2102**, the L1 cache **110** determines that a new address from the first interface is being written to the victim cache in a later pipeline stage, control proceeds to block **2106** to compare the new address to the address of the cache hit. If, at block **2102**, the L1 cache **110**

determines that a new address from the first interface is not being written to the victim cache in a later pipeline stage, then, at block **2104**, the L1 cache **110** determines whether a new address from a second interface has been written to victim cache in a later pipeline stage. For example, the second decoder **860B** can receive VTA_WR_SET1 at the E2 pipeline stage, which can be representative of the scalar interface **502** of FIG. 5 writing an address to the victim storage **218** at the E3 pipeline stage.

(788) If, at block **2104**, the L1 cache **110** determines that a new address from the second interface is not being written to the victim cache in a later pipeline stage, control returns to block **2012** of the example machine readable instructions **2000** of FIG. 20 to output the cache hit address(es) based on the cache hit-miss conversion logic.

(789) If, at block **2104**, the L1 cache **110** determines that a new address from the second interface is being written to the victim cache in a later pipeline stage, then at block **2106**, the L1 cache **110** compares the new address to the address of the cache hit. For example, the first AND gate **864A** can assert a logic one in response to an address of VTAG_WR_SET0 not matching an address of HIT_DP0. In other examples, the third comparator **870A** can compare an address of HIT_DP0 to an address being written to the victim storage **218** by the scalar interface **502**. In yet other examples, the fourth comparator **872A** can compare an address of HIT_DP0 to an address being written to the victim storage **218** by the vector interface **504**.

(790) In response to comparing the new address to address of the cache hit at block **2106**, control proceeds to block **2108** to determine whether the addresses match. If, at block **2108**, the L1 cache **110** determines that the addresses do not match, control returns to block **2012** of the example machine readable instructions **2000** of FIG. 20 to output the cache hit address(es) based on the cache hit-miss conversion logic.

(791) If, at block **2108**, the L1 cache **110** determines that the addresses match, then, at block **2110**, the L1 cache **110** determines whether a cache hit or a cache miss is identified. For example, the first AND gate **864A**, the third comparator **870A**, and/or the fourth comparator **872A** can determine that there is a cache hit of the address of ADP_ADDR_E2_DP0 in the victim storage **218** based on HIT_DP0 including at least one bit value of 1. In other examples, the second AND gate **864B**, the fifth comparator **870B**, and/or the sixth comparator **872B** can determine that there is a cache hit of the address of ADP_ADDR_E2_DP1 in the victim storage **218** based on HIT_DP1 including at least one bit value of 1.

(792) If, at block **2110**, the L1 cache **110** determines that a cache hit is identified, then, at block **2112**, the L1 cache **110** converts the cache hit to a cache miss. For example, the first AND gate **864A** can output a logic low to convert a cache hit to a cache miss. In other examples, the second AND gate **864B** can output a logic low to convert a cache hit to a cache miss. In response to converting the cache hit to a cache miss at block **2112**, control returns to block **2012** of the example machine readable instructions **2000** of FIG. 20 to output the cache hit address(es) based on the cache hit-miss conversion logic.

(793) If, at block **2110**, the L1 cache **110** determines that a cache miss is identified, control proceeds to block **2114** to convert the cache miss to a cache hit in response to the new address matching a write from one(s) of the interface(s) allocating into victim cache. For example, the third comparator **870A** and/or the fourth comparator **872A** can assert a logic one to convert a cache miss to a cache hit in response to ADP_ADDR_E2_DP0 matching an address of a write operation from either DP0 or DP1. In other examples, the fifth comparator **870B** and/or the sixth comparator **872B** can assert a logic one to convert a cache miss to a cache hit in response to ADP_ADDR_E2_DP1 matching an address of a write operation from either DP0 or DP1. In response to converting the cache miss to a cache hit at block **2114**, control returns to block **2012** of the example machine readable instructions **2000** of FIG. 20 to output the cache hit address(es) based on the cache hit-miss conversion logic.

(794) FIG. 22 is an example flowchart representative of example machine readable instructions

2200 that may be executed by the example L1 cache **110** of FIGS. 1-5 to perform data allocation in the main storage **214**, in conjunction with the above description. Although the instructions of FIG. **22** are described in conjunction with the L1 cache **110** of FIGS. 1-5, the instructions may be described in conjunction with any type of storage in any type of cache.

(795) The example main cache controller **222** (FIG. 2) obtains an instruction from the CPU interface **202** (FIG. 2). For example, the CPU interface **202** provides an instruction to the cache controller **220**, the cache controller **220** propagates the instruction to the main cache controller **222**.

(796) The main cache controller **222** determines if the instruction is a read instruction (block **2204**). For example, the main cache controller **222** analyzes the request in the instruction to determine if the main cache controller **222** is to read data from the main storage **214** and provide it to the CPU interface **202** or to perform a different task for the CPU **102**.

(797) If the main cache controller **222** determines the instruction is a read instruction (e.g., block **2204** returns a value YES), the main cache controller **22** determines the address of the read instruction (block **2206**). For example, the main cache controller **222** determines where the data is to be read from in the main storage **214**. In some examples, the main tag RAM access **204** determines the address of the read instruction.

(798) The main cache controller **222** determines if the address of the read instruction matches an address in the tag RAMs **208**, **210**. For example, the cache controller **220** may obtain hit/miss results from the tag RAM access(es) **204**, **206** and determine if the address is available in the main storage **214** and/or victim storage **218**. The main cache controller **222** determines the read instruction is a miss (e.g., block **2208** returns a value NO), the main cache controller **223** identifies the cache line associated with the address (block **2210**). For example, the main cache controller **222** is a direct mapped cache, and the address of the read instruction can only be stored in one location (e.g., at one cache line) of the main storage **214**.

(799) The main cache controller **222** allocates data of the cache line to the victim storage **218** (block **2212**). For example, the main cache controller **222** allocates data from the direct mapped cache line to the victim storage **214**. The main cache controller **222** allocates data regardless of the MESI state of that data. Such an allocation reduces latency of the main cache controller **222** and the overall L1 cache **110** by allocated any line in the main storage **214** to the victim storage **218**.

(800) FIG. **23** is an example flowchart representative of example machine readable instructions **2300** that may be executed by the example L1 cache **110** of FIGS. 1-5 to facilitate a snoop request, in conjunction with the above Section 10. At block **2302**, the snoop address **502** (e.g., the snoop interface) obtains the snoop request from a higher-level data cache (e.g., the L2 data cache **112**).

(Block **2302**). In response, the snoop address **502** issues a read instruction to the tag RAM **210**. (Block **2304**). In examples disclosed herein, the read instruction is issued to the tag RAM **210** to identify whether the victim storage **218** includes the data requested via the snoop address **502**.

(801) At block **2306**, the comparison logic **306c** determines whether the read issued to the tag RAM **210** was a hit. (Block **2306**). In the event the comparison logic **306c** determines the read issued to the tag RAM **210** is not a hit (e.g., the control of block **2306** returns a result of NO), the victim storage **218** generates a snoop response indicating a miss occurred. (Block **2308**).

Additionally, the victim storage **218** transmits the snoop response back to the higher-level data cache (e.g., the L2 data cache **112**). (Block **2310**).

(802) Alternatively, in the event the comparison logic **306** determines the read issued to the tag RAM **210** is a hit (e.g., the control of block **2306** returns a result of YES), the comparison logic **306c** determines the state of the address associated with the read instruction in the MESI RAM **300**. (Block **2312**). At block **2312**, the comparison logic **306** may also store the state of the address as identified responsive to the read instruction in the MESI RAM **300**.

(803) At block **2314**, the example address encoder **326c** generates an address value for use by the victim storage **218** in obtaining the data. (Block **2314**). The address encoder **326c** encodes an address of the tag RAM **210** to a form that is interpretable by the victim storage **218**. For example,

the tag RAM **210** may store 16-bit memory addresses while the victim storage **218** stores 4-bit memory addresses corresponding to the 16-bit memory addresses. Thus, the address encoder **326** may transform the 16-bit memory address into a 4-bit memory address to locate and/or enter the corresponding memory address in the victim storage **218**.

(804) At block **2316**, the example response multiplexer **508** determines whether a data input is obtained from the victim cache store queue **216**. (Block **2316**). In the event the response multiplexer **508** determines no data has been input from the victim cache store queue **216** (e.g., the control of block **2316** returns a result of NO), the response multiplexer **508** outputs the data identified based on the address provided by the address encoder **326c** as the snoop response to the higher-level data cache (e.g., the L2 data cache **112**). (Block **2320**).

(805) Alternatively, in the event the response multiplexer **508** determines data has been input from the victim cache store queue **216** (e.g., the control of block **2316** returns a result of YES), the response multiplexer **508** identifies the updated version of the data as the data to be sent in the snoop response. (Block **2318**). In response to the execution of the control illustrated in block **2318**, the response multiplexer **508** outputs the data identified based on the address provided by the address encoder **326c** as the snoop response to the higher-level data cache (e.g., the L2 data cache **112**). (Block **2320**).

(806) At block **2322**, the snoop address component **506** (e.g., the snoop interface) determines whether an additional snoop request is available. (Block **2322**). In the event the snoop address component **506** (e.g., the snoop interface) determines an additional snoop request is available (e.g., the control of block **2322** returns a result of YES), the process **2300** returns to block **2302**.

Alternatively, in the event the snoop address component **506** (e.g., the snoop interface) determines an additional snoop request is not available (e.g., the control of block **2322** returns a result of NO), the process **2300** stops.

(807) FIGS. **24**, **25**, **26**, **27**, **28**, and **29A**, **29B-1**, and **29B-2** are example flowcharts representative of example machine readable instructions that may be executed by the example L1 cache **110** of FIGS. **1-5** to perform eviction of data in the victim storage **218**, in conjunction with the above Section 11. Although the instructions of FIGS. **24**, **25**, **26**, **27**, **28**, and **29A**, **29B-1**, and **29B-2** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache.

(808) FIG. **24** illustrates an example first operation **2400** of the replacement policy component **308** (FIGS. **3** and **5**) when the first and second data paths (DP0 and DP1) include valid transactions. FIG. **25** illustrates an example second operation **2500** of the replacement policy component **308** when the first and second data paths (DP0 and DP1) include valid transactions. FIG. **26** illustrates an example third operations **2600** of the replacement policy component **308** when the first and second data paths (DP0 and DP1) include valid transactions. FIG. **27** illustrates an example valid-invalid operation **2700** of the replacement policy component **308** when the first data path is a valid transaction and the second data path is an invalid transaction. FIG. **28** illustrates an example invalid-valid operation **2800** of the replacement policy component **308** when the first data path is an invalid transaction and the second data path is a valid transaction. FIG. **29A**, FIG. **29B-1**, and FIG. **29B-2** illustrate an incrementing operation **2900** of the replacement policy component **308** when a data path is allocating data to the victim storage.

(809) Turning to FIG. **24**, the example scalar interface **502** and the example vector interface **504** determine if the first and second data paths are valid transactions (block **2402**). For example, the scalar interface **502** determines if the first data path DP0 is accessing (e.g., requesting a read or write operation) the victim storage **218** and the vector interface **504** determines if the second data path DP1 is accessing the victim storage **218**. When the interfaces **502**, **504** determine the first and second data paths are valid transactions (e.g., block **2402** returns a value YES), the replacement policy component **308** obtains results from the hit-miss comparison logic **306a**, **306b** (block **2404**). For example, the replacement policy component **308** obtains results indicating whether the first

data path access and the second data path access have matching addresses in the tag RAM **210** or do not.

(810) The example replacement policy component **308** determines if the results indicate that both of the transactions of the first data path and the second data path are hits (block **2406**). When the replacement policy component **308** determines DP0 and DP1 are both hits (e.g., block **2406** returns a value yes), the least recently used value Y remains constant (block **2408**). For example, since neither the first data path nor the second data path needs to evict data, the LRU value does not need to change.

(811) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2410**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2410** returns a value YES) then control returns to block **2402**. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2410** returns a value NO) then the first operation **2400** ends.

(812) When the replacement policy component **308** determines the results do not indicate that both of the transactions of the first data path and the second data path are hits (e.g., block **2406** returns a value NO), the replacement policy component **308** determines if the results indicate that both of the transactions of the first data path and the second data path are misses (block **2412**). For example, the replacement policy component **308** determines if both results from the first hit-miss comparison logic **306a** and the second hit-miss comparison logic **306b** indicate neither of the accesses matched the addresses in the tag RAM **210**.

(813) If DP0 and DP1 missed (e.g., block **2412** returns a value YES), then the replacement policy component **308** points the first data path to the LRU way (Y) (block **2414**). The replacement policy component **308** points the second data path to the next LRU way (Y+1) (block **2416**). For example, the victim storage **218** includes n number of ways, each way has a location (e.g., slot 1, slot 2, slot n), each way is mapped to an address, and each way includes data. The replacement policy component **308** initializes a value Y to be equal to the least recently used way in the victim cache. For example, the LRU way is slot 2, thus Y is equal to slot 2. When the replacement policy component **308** points the first data path to the LRU way (block **2414**), the replacement policy component **308** is assigning the location of Y in the victim storage **218** to DP0 for eviction. Similarly, when the replacement policy component **308** points the second data path to the next LRU way (block **2416**), the replacement policy component **308** is assigning the location of Y+1 in the victim storage **218** to DP1 for eviction.

(814) The example replacement policy component **308** provides the pointer values to the example multiplexers **330a**, **330b** (block **2418**). For example, the replacement policy component **308** provides a location (Y) of the way that is to be evicted by DP0 from the victim storage **218** to the multiplexer **330a** and a location (Y+1) of the way that is to be evicted by DP1 to the multiplexer **330b**. In some examples, the selecting input of the multiplexer **330a** and **330b** selects the replacement policy component input, the address read **332a** and **332b** reads the input of the replacement policy component **308** and evicts the ways indicated by location Y and location Y+1.

(815) The example replacement policy component **308** increments Y by two (block **2420**). For example, the replacement policy component **308** assigns Y to a location in the victim storage **218** that is two slots away from the current value of Y (e.g., if Y=slot 1, then Y=slot 3 after incrementing). In this manner, during the next clock cycle, the replacement policy component **308** is provided with an updated and appropriate Y value. Alternatively and/or additionally, the replacement policy component **308** increments indicators for the first and second data paths by two.

(816) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2422**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous

accesses (e.g., block **2422** returns a value YES) then control returns to block **2402**. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2420** returns a value NO) then the first operation **2400** ends.

(817) Turning to FIG. **25**, when the replacement policy component **308** determines the results indicate that the first data path and the second data path are not both misses (e.g., block **2412** returns a value NO), the example replacement policy component **308** determines if the results indicate that the first data path is a hit and the second data path is a miss (block **2502**).

(818) If the example replacement policy component **308** determines that the first data path is a hit and the second data path is a miss (e.g., block **2502** returns a value YES), then the replacement policy component determines the location in the victim storage **218** of the hit way (DP0 Way) (block **2504**). For example, the replacement policy component **308** analyzes the address of DP0 and identifies the location in the victim storage **218** that includes that address. In some examples, the replacement policy component **308** may include an updated list of the elements in the victim storage **218**. In other examples, the replacement policy component **308** retrieves and/or obtains information from the tag RAM **210** regarding the locations of the addresses stored in the victim storage **218**.

(819) The example replacement policy component **308** determines if the hit way (DP0 Hit Way) matches the location of the next LRU value (Y+1) (block **2506**). For example, the replacement policy component **308** may compare the location of the hit way containing the address of DP0 to the location value assigned to Y+1. If the replacement policy component **308** determines that the locations are match (e.g., block **2506** returns a value YES), then the replacement policy component **308** switches the assignment of the next LRU value and the LRU value (block **2508**). For example, the second data path DP1 pointer is to be assigned to the LRU value (e.g., location Y) instead of the next LRU value (e.g., location Y+1). The replacement policy component **308** switches the assignment to avoid the second data path DP1 evicting the DP0 Hit Way. In some examples, the replacement policy component **308** decrements an indicator to indicate the LRU way of the victim storage **214** to be evicted by the second data path DP1.

(820) The replacement policy component **308** points the second data path to the LRU way (Y) (block **2510**). For example, the replacement policy component **308** assigns the value of Y (e.g., the location of the LRU way) to the second data path DP1 for eviction.

(821) The example replacement policy component **308** provides the pointer values to the multiplexer(s) **330a**, **330b** (block **2512**). For example, the replacement policy component **308** provides a location (Y) of the way that is to be evicted by DP1 from the victim storage **218** to the multiplexer **330b**. In some examples, when the hit way does not match the location of the next LRU value (e.g., block **2506** returns a value NO), the replacement policy component **308** provides the pointer value Y+1 and the location of the hit way to the multiplexer(s) **330a**, **330b**. For example, the original assignment of the next LRU value to the second data path DP1 remains the same.

(822) The example replacement policy component **308** increments Y based on eviction (block **2514**). For example, if the assignments of LRU values to data paths were switched (e.g., DP1 pointer points to the LRU value Y), then the replacement policy component **308** increments Y by one. Otherwise, the replacement policy component **308** increments Y by two. In this manner, during the next clock cycle, the replacement policy component **308** is provided with an updated Y value and Y+1 value. Alternatively and/or additionally, the replacement policy component **308** increments indicators for the first and second data paths based on eviction.

(823) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2516**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2516** returns a value YES) then control returns to block **2402** of FIG. **24**. If the

replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2516** returns a value NO) then the second operation **2500** ends.

(824) Turning to FIG. **26**, when the replacement policy component **308** determines the results indicate that the first data path is not a hit and the second data path is not a miss (e.g., block **2602** returns a value NO), the example replacement policy component **308** determines that the results indicate that the first data path is a miss and the second data path is a hit (block **2602**).

(825) The example replacement policy component **308** determines the location in the victim storage **218** of the hit way (DP1 Way) (block **2604**). For example, the replacement policy component **308** analyzes the address of DP1 and identifies the location in the victim storage **218** that includes that address.

(826) The example replacement policy component **308** determines if the hit way (DP1 Way) matches the location of the LRU value (Y) (block **2606**). For example, the replacement policy component **308** may compare the location of the hit way containing the address of DP1 to the location value assigned to Y. If the replacement policy component **308** determines that the locations match (e.g., block **2606** returns a value YES), then the replacement policy component **308** switches the assignment of the LRU value and the next LRU value (block **2608**). For example, first data path DP0 pointer is to be assigned to the next LRU value (e.g., location Y+1) instead of the LRU value (e.g., location Y). The replacement policy component **308** switches the assignment to avoid the first data path DP0 evicting the DP1 Hit Way. In some examples, the replacement policy component **308** increments an indicator to indicate the next LRU way in the victim storage **214** to be evicted by the first data path DP0.

(827) The replacement policy component **308** points the first data path to the next LRU value (Y+1) (block **2610**). For example, the replacement policy component **308** assigns the value of Y+1 (e.g., the location of the next LRU way) to the first data path DP0 for eviction.

(828) The example replacement policy component **308** provides the pointer values to the multiplexer(s) **330a**, **330b** (block **2612**). For example, the replacement policy component **308** provides a location (Y+1) of the way that is to be evicted, by DP0, from the victim storage **218** to the multiplexer **330a**. In some examples, when the hit way does not match the location of the LRU value (e.g., block **2506** returns a value NO), the replacement policy component **308** provides the pointer value Y and the location of the hit way to the multiplexer(s) **330a**, **330b**. For example, the original assignment of the LRU value to the first data path DP0 remains the same.

(829) The example replacement policy component **308** increments Y based on eviction (block **2614**). For example, if the assignments of LRU values to data paths were switched (e.g., DP0 pointer points to the next LRU value Y+1), then the replacement policy component **308** increments Y by two. Otherwise, the replacement policy component **308** increments Y by one. In this manner, during the next clock cycle, the replacement policy component **308** is provided with an updated Y value. Alternatively and/or additionally, the replacement policy component **308** increments indicators for the first and second data paths based on eviction.

(830) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2616**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2616** returns a value YES) then control returns to block **2402** of FIG. **24**. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2616** returns a value NO) then the third operation **2600** ends.

(831) Turning to FIG. **27**, when the example scalar interface **502** and the example vector interface **504** determine the first and second data paths are not valid transactions (e.g., block **2402** returns a value NO), the example scalar interface **502** and the example vector interface **504** determine if the first data path is valid and the second data path is invalid (block **2702**). For example, the scalar

interface **502** determines if the first data path DP0 is accessing (e.g., requesting a read or write operation) the victim storage **218** and the vector interface **504** determines if the second data path DP1 is not attempting to access the victim storage **218**.

(832) When the interfaces **502**, **504** determine the first data path is a valid transaction and the second data path is an invalid transaction (e.g., block **2702** returns a value YES), the replacement policy component **308** obtains results from the hit-miss comparison logic **306a** (block **2704**). For example, the replacement policy component **308** obtains a result indicating whether the first data path access has a matching addresses in the tag RAM **210** or does not have a matching address in the tag RAM **210**. The example replacement policy component **308** determines if the results indicate that first data path is a hit (block **2706**). If the replacement policy component **308** determines the address of the first data path DP0 hits an address in the tag RAM **210** (e.g., block **2706** returns a value YES), the least recently used value Y remains constant (block **2708**). For example, since the first data path does not need to evict data, the LRU value does not need to change.

(833) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2710**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2710** returns a value YES) then control returns to block **2402** of FIG. 24. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2710** returns a value NO) then the first operation **2700** ends.

(834) When the example replacement policy component **308** determines that the results do not indicate that first data path is a hit (e.g., block **2706** returns a value NO), then the first data path is a miss (block **2712**).

(835) The example replacement policy component **308** points the first data path to the LRU Way (Y) (block **2714**). For example, the replacement policy component **308** assigns the location of Y in the victim storage **218** to DP0 for eviction.

(836) The example replacement policy component **308** provides the pointer value to the first multiplexer **330a** (block **2716**). For example, the replacement policy component **308** provides the location the LRU way to the first multiplexer **330a** for eviction of that way.

(837) The example replacement policy component **308** increments Y (block **2718**). For example, the replacement policy component **308** updates the LRU way to the next location (e.g., Y+1) in the victim storage **218**. Alternatively and/or additionally, the replacement policy component **308** increments indicators for the first and second data paths.

(838) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2720**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2720** returns a value YES) then control returns to block **2402** of FIG. 24. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2720** returns a value NO) then the first operation **2700** ends.

(839) Turning to FIG. 28, when the example scalar interface **502** and the example vector interface **504** determine the first data path is not a valid transaction and the second data path is a valid transactions (e.g., block **2702** returns a value NO), the example scalar interface **502** and the example vector interface **504** determine the first data path is invalid and the second data path is valid (block **2802**). For example, the scalar interface **502** determines that the first data path DP0 is not accessing (e.g., requesting a read or write operation) the victim storage **218** and the vector interface **504** determines if the second data path DP1 is accessing the victim storage **218**.

(840) The replacement policy component **308** obtains results from the hit-miss comparison logic **306b** (block **2804**). For example, the replacement policy component **308** obtains a result indicating

whether the second data path access has a matching addresses in the tag RAM **210** or does not have a matching address in the tag RAM **210**.

(841) The example replacement policy component **308** determines if the results indicate that second data path is a hit (block **2806**). If the replacement policy component **308** determines the address of the second data path DP1 hits an address in the tag RAM **210** (e.g., block **2806** returns a value YES), the least recently used value Y remains constant (block **2808**). For example, since the second data path does not need to evict data, the LRU value does not need to change.

(842) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2810**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2810** returns a value YES) then control returns to block **2402** of FIG. 24. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2810** returns a value NO) then the first operation **2700** ends.

(843) When the example replacement policy component **308** determines that the results do not indicate that second data path is a hit (e.g., block **2806** returns a value NO), then the second data path is a miss (block **2812**).

(844) The example replacement policy component **308** points the second data path to the LRU Way (Y) (block **2814**). For example, the replacement policy component **308** assigns the location of Y in the victim storage **218** to DP1 for eviction.

(845) The example replacement policy component **308** provides the pointer value to the second multiplexer **330b** (block **2816**). For example, the replacement policy component **308** provides the location the LRU way to the second multiplexer **330b** for eviction of that way.

(846) The example replacement policy component **308** increments Y (block **2818**). For example, the replacement policy component **308** updates the LRU way to the next location (e.g., Y+1) in the victim storage **218**. Alternatively and/or additionally, the replacement policy component **308** increments indicators for the first and second data paths.

(847) The example replacement policy component **308** determines if a transaction on a new clock cycle has been received (block **2820**). For example, if the replacement policy component **308** obtains hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2820** returns a value YES) then control returns to block **2402** of FIG. 24. If the replacement policy component **308** does not obtain hit-miss results corresponding to different transactions (e.g., accesses) than the previous accesses (e.g., block **2820** returns a value NO) then the first operation **2800** ends.

(848) The machine readable instructions **2400**, **2500**, **2600**, **2700**, and **2800** correspond to the first table **602** of FIG. 6.

(849) FIG. 29A, FIG. 29B-1, and FIG. 29B-2 are example flowcharts representative of example machine readable instructions **2900** that may be executed by the L1 cache **110** of FIGS. 1-5 to perform LRU incrementing in the victim storage **214** based on the allocation status of a data path, in conjunction with the above description.

(850) The machine readable instructions **2900** begin at block **2902**, at which the replacement policy component **308** initializes the first data path allocate pointer to equal location Y. For example, the replacement policy component **308** assigns a portion of the victim storage **218** not recently used by the CPU **102**, having location Y, to the LRU value. In such an example, when the first data path DP0 is to allocate, the victim storage **218** evicts data from the LRU value (e.g., location Y). The replacement policy component **308** initializes the second data path allocate pointer to equal location Y+1 (block **2904**). For example, the replacement policy component **308** assigns a portion of the victim storage **218** not recently used by the CPU **102**, having location Y+1, to the next LRU value. In such an example, when the second data path DP1 is to allocate, the victim storage **218** evicts data from the next LRU value (e.g., location Y+1).

(851) The replacement policy component **308** determines the first and second data paths are valid transactions (block **2906**). For example, the CPU **102** provided instructions on both data paths. (852) The replacement policy component **308** determines if the first data path DP0 is a hit in the victim storage **218** and the second data path DP1 is to allocate (block **2908**). If the first data path DP0 hits the victim storage **218** and the second data path DP1 is allocating to the victim storage **218** (e.g., block **2908** returns a value YES), the replacement policy component **308** determines the location in the victim storage **218** of the of the hit location (DP0 Hit Way) (block **2910**). For example, the replacement policy component **308** identifies where the address of the DP0 is in the victim storage **218**.

(853) The replacement policy component **308** determines if the hit location is equal to the location of the second data path allocate pointer (Y+1) (block **2912**). For example, the replacement policy component **308** determines if the address of DP0 matches the location of Y+1. If the locations do match (e.g., block **2912** returns a value YES), the replacement policy component **308** updates the second data path allocate pointer to equal location Y (block **2914**). For example, the replacement policy component **308** switches the assignment of the second data path allocate pointer from Y+1 to Y to avoid evicting data requested on the DP0 instruction.

(854) The cache controller **220** performs the first transaction and the second transaction (block **2916**). For example, the cache controller **220** reads/writes data of DP0 at location Y+1 and evicts data from location Y in the victim storage **218**.

(855) The replacement policy component **308** increments the first data path allocate pointer by one (block **2918**). For example, since the cache controller **220** evicted data from location Y and not Y+1, the replacement policy component **308** only needs to update the LRU value to the next LRU value (Y+1).

(856) If the hit location and the second data path allocate pointer do not match (e.g., block **2912** returns a value NO), the cache controller **220** performs the first transaction and the second transaction (block **2920**). For example, the replacement policy component **308** determines that Y+1 includes data that is available to evict and thus, the second data path allocate pointer can evict data from that location while the first data path DP0 reads/writes data from the hit location.

(857) The replacement policy component **308** increments the first data path allocate pointer by two (block **2922**). For example, since the cache controller **220** evicts data from the location Y+1, the replacement policy component **308** updates the LRU value Y to a location after the evicted location (e.g., LRU value=Y+2 and the next LRU value=Y+3).

(858) If the replacement policy component **308** determines the condition of block **2908** is not true (e.g., block **2908** returns a value NO), the replacement policy component **308** determines if the first data path is to allocate and the second data path hits (block **2924**). For example, the replacement policy component **308** determines if the second data path hits a location in the victim storage **218** and if the main storage **214** is allocating data on the first data path DP0.

(859) If the replacement policy component **308** determines the first data path is to allocate and the second data path is a hit (e.g., block **2924** returns a value YES), the replacement policy component **308** determines the location in the victim storage **218** of the hit location (DP2 Way) (block **2926**). For example, the replacement policy component **308** determines where the second data path is reading/writing data from in the victim storage **218**.

(860) Turning to FIG. **29B-1**, the replacement policy component **308** determines if the hit location is equal to the location Y (block **2928**). For example, the replacement policy component **308** determines if the first data path allocate pointer points to the same location storing the hit data.

(861) If the replacement policy component **308** determines the locations match (e.g., block **2928** returns a value YES), the replacement policy component **308** updates the first data path allocate pointer to equal location Y+1 (block **2930**). For example, the replacement policy component **308** switches the assignments of the LRU value and the next LRU value to avoid the first data path evicting the hit data from the victim storage **218**.

(862) The cache controller **220** performs the first transaction and the second transaction (block **2932**). For example, the cache controller **220** reads/writes data from the location Y and evicts data of the location Y+1. The replacement policy component **308** increments the pointer Y by two locations (block **2934**). For example, the replacement policy component **308** updates the LRU location to a location after the most recently evicted location (e.g., in this example, the most recently evicted location is Y+1, therefore the LRU value is incremented by two to equal Y+2).

(863) If the replacement policy component **308** determines the hit location does not match the location of the originally assigned first data path pointer (Y) (e.g., block **2928** returns a value NO), the cache controller **220** performs the first transaction and the second transaction (block **2936**). For example, the replacement policy component **308** determines that Y includes data that is available to evict and thus, the first data path allocate pointer can evict data from that location while the second data path DP1 reads/writes data from the hit location.

(864) The replacement policy component **308** increments the first data path allocate pointer by one location (block **2938**). For example, since the cache controller **220** evicts data from the location Y, the replacement policy component **308** updates the LRU value Y to a location after the evicted location. In this manner, the replacement policy component **308** includes an updated LRU value and an updated next LRU value during the next clock cycle.

(865) If the condition of block **2924** of FIG. **29A** is not true (e.g., block **2924** returns a value NO when the first data path is not to allocate when the second data path is to hit), then control moves to FIG. **29B-2** where the replacement policy component **308** determines both data paths are to allocate (block **2940**). For example, if two read-misses occur, the main storage **214** allocates two lines from the main storage **214** to the victim storage **218**.

(866) In this manner, the cache controller **220** performs the first transaction and the second transaction (block **2942**). For example, the cache controller **220** evicts data from the LRU location (Y) utilizing the first data path DP0 and evicts data from the next LRU location (Y+1) utilizing the second data path DP1.

(867) The replacement policy component **308** increments the first data path allocate pointer by two (block **2944**). For example, the replacement policy component **308** increments the location Y by two locations since data was evicted from Y+1. In some examples, when the LRU value is incremented by a value, the next LRU value is incremented simultaneously by the same value. Therefore, the first data path allocate pointer and the second data path allocate pointer always point to an updated and accurate eviction location.

(868) The machine readable instructions **2900** correspond to the second table **604** of FIG. **6**.

(869) FIG. **30** is an example flowchart representative of example machine readable instructions **3000** that may be executed by the example L1 cache **110** of FIGS. **1-5** to execute arbitration logic to perform a read, modify, or write operation as described above. Although the instructions of FIG. **30** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **30** are described in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.). However, the instructions of FIG. **30** can likewise be used in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.).

(870) The machine readable instructions **3000** of FIG. **30** begin at block **3002**, at which the L1 cache **110** obtains store instruction(s) from interface(s) coupled to hardware. For example, the address processing components **302a-c** of FIGS. **3A-3D** can obtain the first store instruction **1018** from the scalar interface, the second store instruction **1020** from the memory interface, and/or the third store instruction **1022** from the vector interface. Alternatively, the address processing components **302a-c** of FIGS. **3A-3D** can obtain the first store instruction **1018b** from the scalar interface, the second store instruction **1020b** from the memory interface, and/or the third store instruction **1022b** from the vector interface of FIG. **10B**.

(871) At block **3004**, the L1 cache **110** generates transaction data based on the store instruction(s). For example, the address processing components **302a-c** and/or the bank processing logic **303** of FIGS. **3A-3D** can generate the first transaction data, the second transaction data, and/or the third transaction data of FIG. **10A**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** of FIGS. **3A-3D** can generate the first transaction data, the second transaction data, and/or the third transaction data of FIG. **10B**. An example process that may be used to implement block **3004** is described below in connection with FIG. **31**.

(872) At block **3006**, the L1 cache **110** determines whether read operation(s) is/are identified based on the transaction data. For example, the address processing components **302a-c** can determine that at least one of the first store instruction **1018**, the second store instruction **1020**, or the third store instruction **1022** includes a request to have a read operation serviced (e.g., a value of `RD_BANK_REQ[i]` is indicative of a read request, a logic high signal for `RD_BANK_REQ[i]`, etc.). In such examples, the read operation request can be determined based on the R/W data included in the store instructions **1018**, **1020**, **1022**. Alternatively, the address processing components **302a-c** can determine that at least one of the first store instruction **1018b**, the second store instruction **1020b**, or the third store instruction **1022b** includes a request to have a read operation serviced (e.g., a value of `RD_BANK_REQ[i]` is indicative of a read request, a logic high signal for `RD_BANK_REQ[i]`, etc.). In such examples, the read operation request can be determined based on the R/W data included in the store instructions **1018b**, **1020b**, **1022b** of FIG. **10B**.

(873) If, at block **3006**, the L1 cache **110** determines that there are no read operations identified based on the transaction data, control proceeds to block **3014** to invoke second arbitration logic to write the data to the store queue. If, at block **3006**, the L1 cache **110** determines that there is at least one read operation identified based on the transaction data, then, at block **3008**, the L1 cache **110** invokes first arbitration logic to locate data for the read operation(s) in at least one of a store queue or storage. For example, the address processing components **302a-c** can invoke the first arbitration logic **1008** to locate data for the read operation(s) in at least one of the main cache store queue **212** or the main storage **214**. Alternatively, the address processing components **302a-c** can invoke the first arbitration logic **1008b** to locate data for the read operation(s) in at least one of the victim cache store queue **216** or the victim storage **218**. An example process that may be executed to implement block **3008** is described below in connection with FIG. **20**.

(874) At block **3010**, the L1 cache **110** identifies the most recent version of the located data. For example, the L1 cache **110** can compare a first version of the requested data from the main cache store queue **212** to a second version of the requested data from the main storage **214** and determine that the first version is more recent than the second version based on the comparison. Alternatively, the L1 cache **110** can compare a first version of the requested data from the victim cache store queue **216** to a second version of the requested data from the victim storage **218** and determine that the first version is more recent than the second version based on the comparison.

(875) At block **3012**, the L1 cache **110** delivers the most recent version of the located data to store queue to execute a modify operation on the read and write data. For example, the main cache store queue **212** can deliver and/or otherwise transmit the first version of the requested data to the main cache store queue **212** to execute a modify operation on the requested data and the data to be written. Alternatively, the victim cache store queue **216** can deliver and/or otherwise transmit the first version of the requested data to the victim cache store queue **216** to execute a modify operation on the requested data and the data to be written.

(876) At block **3014**, the L1 cache **110** invokes the second arbitration logic to write the data to the store queue or the storage. For example, the first arbitration logic **1008** can transmit an instruction to the second arbitration logic **1010** to write the WDATA or portion(s) thereof to at least one of the main cache store queue **212** or the main storage **214**. Alternatively, the first arbitration logic **1008b** can transmit an instruction to the second arbitration logic **1010b** to write the WDATA or portion(s)

thereof to at least one of the victim cache store queue **216** or the victim storage **218** of FIG. **10B**. In response to invoking the second arbitration logic to write the data to the store queue or the storage at block **3014**, the example machine readable instructions **3000** of FIG. **30** conclude.

(877) FIG. **31** is an example flowchart representative of example machine readable instructions **3100** that may be executed by the example L1 cache **110** of FIGS. **1-5** to generate transaction data based on store instruction(s) as described above. The flowchart of FIG. **31** can be an example implementation of the machine readable instructions **3004** of FIG. **30**. Although the instructions of FIG. **31** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **31** are described in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.). However, the instructions of FIG. **31** can likewise be used in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.).

(878) The machine readable instructions **3100** begin at block **3102**, at which the L1 cache **110** extracts write data from the store instruction(s) based on a number of data storage banks. For example, the address processing components **302a-c** and/or the bank processing logic **303** of FIGS. **3A-3D** can extract the WDATA from the store instructions **1018**, **1020**, **1022** based on a quantity of data banks that the main cache store queue **212** and/or the main storage **214** are broken up into. In such examples, in response to the main cache store queue **212** having 16 data banks, the cache line can be 64 bits and, thus, WDATA can be extracted in 64 bit chunks. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** of FIGS. **3A-3D** can extract the WDATA from the store instructions **1018b**, **1020b**, **1022b** of FIG. **10B** based on a quantity of data banks that the victim cache store queue **216** and/or the victim storage **218** are broken up into. In such examples, in response to the victim cache store queue **216** having 16 data banks, the cache line can be 64 bits and, thus, WDATA can be extracted in 64 bit chunks.

(879) At block **3104**, the L1 cache **110** rotates the write data to generate rotated write data. For example, the address processing components **302a-c** can rotate the CPU-Write Data of the first store instruction **902** of FIGS. **3A-3D** to generate the rotated data **904** of FIG. **9**.

(880) At block **3106**, the L1 cache **110** determines byte enable data based on the store instruction(s). For example, the address processing components **302a-c** and/or the bank processing logic **303** can determine the BYTEN/BANK[i] data of FIG. **10A** based on the BYTEN data included in the store instructions **1018**, **1020**, **1022**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can determine the BYTEN/BANK[i] data of FIG. **10B** based on the BYTEN data included in the store instructions **1018b**, **1020b**, **1022b**.

(881) At block **3108**, the L1 cache **110** determines a data access operation data size based on the store instruction(s). For example, the address processing components **302a-c** and/or the bank processing logic **303** can determine the data size of data to be read, written, and/or modified based on the SIZE data included in the store instructions **1018**, **1020**, **1022**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can determine the data size of data to be read, written, and/or modified based on the SIZE data included in the store instructions **1018b**, **1020b**, **1022b**.

(882) At block **3110**, the L1 cache **110** determines a data storage address based on the store instruction(s). For example, the address processing components **302a-c** and/or the bank processing logic **303** can determine the MS_ADDR[i] of a corresponding bank of the main cache store queue **212** and/or the STQ_ADDR[i] address of a corresponding bank of the main storage **214** based on the ADDR data included in the store instructions **1018**, **1020**, **1022**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can determine the VS_ADDR[i] of a corresponding bank of the victim cache store queue **216** and/or the STQ_V_ADDR[i] address of a corresponding bank of the victim storage **218** based on the ADDR data included in the store instructions **1018b**, **1020b**, **1022b**.

(883) At block **3112**, the L1 cache **110** maps the data access operation data size and the data storage address to a first quantity of data banks to read from. For example, the address processing components **302a-c** and/or the bank processing logic **303** can map the data access operation size and the data storage address to zero or more banks of the main cache store queue **212**, zero or more banks of the main storage **214**, etc., to generate RD_BANK_REQ[i] of FIG. **10A**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can map the data access operation size and the data storage address to zero or more banks of the victim cache store queue **216**, zero or more banks of the victim storage **218**, etc., to generate RD_BANK_REQ[i] of FIG. **10B**.

(884) At block **3114**, the L1 cache **110** maps the data access operation data size and the data storage address to a second quantity of data banks to write to. For example, the address processing components **302a-c** and/or the bank processing logic **303** can map the data access operation size and the data storage address to zero or more banks of the main cache store queue **212**, zero or more banks of the main storage **214**, etc., to generate WR_BANK_REQ[i] of FIG. **10A**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can map the data access operation size and the data storage address to zero or more banks of the victim cache store queue **216**, zero or more banks of the victim storage **218**, etc., to generate WR_BANK_REQ[i] of FIG. **10B**.

(885) At block **3116**, the L1 cache **110** generates transaction data based on at least one of the first quantity, the second quantity, the byte enable data, or the write data. For example, the address processing components **302a-c** and/or the bank processing logic **303** can generate the first transaction data (TRANSACTION_DP0[i]), the second transaction data (TRANSACTION_DMA[i]), and the third transaction data (TRANSACTION_DP1[i]) of FIG. **10A**. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can generate the first transaction data (TRANSACTION_DP0[i]), the second transaction data (TRANSACTION_DMA[i]), and the third transaction data (TRANSACTION_DP1[i]) of FIG. **10B**. In response to generating the transaction data based on at least one of the first quantity, the second quantity, the byte enable data, or the write data at block **3116**, control returns to block **3006** of the machine readable instructions **3000** of FIG. **30** to determine whether read operation(s) is/are identified based on the transaction data.

(886) FIG. **32** is an example flowchart representative of example machine readable instructions **3200** that may be executed by the example L1 cache **110** of FIGS. **1-5** to invoke first arbitration logic to locate data for read operation(s) in at least one of a store queue or storage as described above. The flowchart of FIG. **32** can be an example implementation of the machine readable instructions **3008** of FIG. **30**. Although the instructions of FIG. **32** are described in conjunction with the L1 cache **110** of FIGS. **1-5**, the instructions may be described in conjunction with any type of storage in any type of cache. The instructions of FIG. **32** are described in conjunction with the main half of the L1 cache **110** (e.g., the main cache store queue **212**, the main storage **214**, etc.). However, the instructions of FIG. **32** can likewise be used in conjunction with the victim side of the L1 cache **110** (e.g., the victim cache store queue **216**, the victim storage **218**, etc.).

(887) The machine readable instructions **3200** begin at block **3202**, at which the L1 cache **110** selects a data storage bank of interest to process. For example, the address processing components **302a-c** and/or the bank processing logic **303** of FIGS. **3A-3D** can select the first bank **1002** of FIG. **10A** to process. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** of FIGS. **3A-3D** can select the first bank **1002b** of FIG. **10B** to process.

(888) At block **3204**, the L1 cache **110** compares selected data storage bank to data storage banks included in read bank request(s) from interface(s). For example, the first arbitration logic **1008** can compare the bank(s) identified in respective one(s) of RD_BANK_REQ[i] from the scalar interface, the memory interface, and the vector interface to the first bank **1002** (e.g., STQ[0], MS[i], etc.). Alternatively, the first arbitration logic **1008b** can compare the bank(s) identified in

respective one(s) of RD_BANK_REQ[i] from the scalar interface, the memory interface, and the vector interface to the first bank **1002b** (e.g., STQ_V[0], VS[i], etc.).

(889) At block **3206**, the L1 cache **110** determines whether at least one interface requests access to the selected data storage. If, at block **3206**, the L1 cache **110** determines that none of the interfaces request access to the selected data bank, control proceeds to block **3208** to determine that the selected data storage bank is not used for read operation(s). In response to determining that the selected data storage bank is not used for read operation(s) at block **3208**, control returns to block **3202** to select another data storage bank of interest to process.

(890) If, at block **3206**, the L1 cache **110** determines that at least one interface requests access to the selected data bank, control proceeds to block **3210** to determine whether more than one interface requests access to the selected data storage bank. If, at block **3210**, the L1 cache **110** determines that only one interface requests access to the selected data storage bank, control proceeds to block **3212** to invoke first arbitration logic to assign the selected data storage bank to the requesting interface. In response to invoking the first arbitration logic to assign the selected data storage bank to the requesting interface at block **3212**, control returns to block **3202** to select another data storage bank of interest to process.

(891) If, at block **3210**, the L1 cache **110** determines that more than one interface requests access to the selected data storage bank, control proceeds to block **3214** to invoke the first arbitration logic to assign the selected data storage bank to the requesting interface requiring a read operation. For example, the first arbitration logic **1008** can assign the first bank **1002** to the one of the interfaces requiring a read operation as read operations are prioritized over write operations. Alternatively, the first arbitration logic **1008b** can assign the first bank **1002b** to the one of the interfaces requiring a read operation as read operations are prioritized over write operations.

(892) At block **3216**, the L1 cache **110** invokes the first arbitration logic to inform second arbitration logic that the requesting interface requiring a write operation is not assigned the selected data storage bank. For example, if scalar data path is requesting a write operation and a read operation and the scalar data path is not assigned the first data bank **1002** for the read operation, the first arbitration logic **1008** can instruct the second arbitration logic **1010** to not assign the scalar data path the first data bank **1002** and, thus, stall and/or otherwise prevent execution of the write operation since the corresponding read operation is not to be completed during the clock cycle. Alternatively, if scalar data path is requesting a write operation and a read operation and the scalar data path is not assigned the first data bank **1002b** for the read operation, the first arbitration logic **1008** can instruct the second arbitration logic **1010b** to not assign the scalar data path the first data bank **1002b** and, thus, stall and/or otherwise prevent execution of the write operation since the corresponding read operation is not to be completed during the clock cycle.

(893) At block **3218**, the L1 cache **110** determines whether to select another data storage bank of interest to process. For example, the address processing components **302a-c** and/or the bank processing logic **303** can determine to select a second bank of the main cache store queue **212** and the main storage **214** to process. Alternatively, the address processing components **302a-c** and/or the bank processing logic **303** can determine to select a second bank of the victim cache store queue **216** and the victim storage **218** to process. If, at block **3218**, the L1 cache **110** determines to select another data storage bank of interest to process, control returns to block **3202** to select another data storage bank of interest to process. If, at block **3218**, the L1 cache **110** determines not to select another data storage bank of interest to process, control returns to block **3010** of the machine readable instructions **3000** of FIG. **30** to identify the most recent version of the located data.

(894) FIG. **33** is an example flowchart representative of example machine readable instructions **3300** that may be executed by the example L1 cache **110** of FIGS. **1-5** to facilitate a read-modify-write operation in the victim storage **216**, in conjunction with the above Section 13. In FIG. **33**, the victim cache store queue **216** obtains a write instruction transmitted by the CPU **102** (e.g.,

transmitted through the cache controller **220**) indicating byte(s) of a word, or an entire word, to be re-written. (Block **3302**). In some examples disclosed herein, the write port **1126** may obtain the write instruction transmitted by the CPU **102**.

(895) At block **3304**, the victim cache store queue **216** transmits the value of the portion of the word to be rewritten to the latch **1102b**. (Block **3304**). In some examples, the latch **1102b** transmits the value of the portion of the word to be rewritten to the latch **1102c**.

(896) At block **3306**, the victim cache store queue **216** stores the address value associated with the location of the portion of the word to be rewritten in the pending store address data store **1116**. (Block **3306**). Additionally, the tag ram **210** transmits a read instruction (e.g., a read request) of the entire currently stored word to the victim storage **218**. (Block **3308**). At block **3310**, the victim cache store queue **216** determines whether there has been a subsequent clock cycle of the CPU **102**, or the cache controller **220**. (Block **3310**). In some examples disclosed herein, the latch **1102c** determines whether there has been a subsequent clock cycle of the CPU **102**, or the cache controller **220**. In response to determining that there has not been a subsequent clock cycle of the CPU **102**, or the cache controller **220**, (e.g., the control of block **3310** returns a result of NO), the process waits.

(897) Alternatively, in response to determining that there has been a subsequent cycle of the CPU **102**, or the cache controller **220**, (e.g., the control of block **3310** returns a result of YES), the read-modify-write merge component **1108** obtains the value of the portion of the word (e.g., the byte) stored in the latch **1102c**. (Block **3312**). Additionally, the read-modify-write merge component **1108** obtains the entire currently stored word transmitted by the ECC logic **312**. (Block **3314**). In this manner, the read-modify-write merge **1108** identifies the address of the byte in the currently stored word to be updated. Once the read-modify-write merge component **1108** identifies and/or otherwise obtains (a) the value (e.g., byte value, bit value, etc.) of the portion of the currently stored word to be updated from the latch **1102c** and the (b) currently stored word from the ECC logic **312**, the read-modify-write merge component **1108** writes (e.g., replaces) the portion of the currently stored word with the value of the portion of the currently stored word obtained from the latch **1102c**. (Block **3316**). For example, the read-modify-write merge component **1108** writes the value of the portion of the word to an address value corresponding to the portion of the word in the word.

(898) At block **3318**, the victim cache store queue **216** generates error detection code based on the word, the error detection code to be stored with the word. (Block **3318**). In some examples disclosed herein the ECC generator **1112** generates error detection code based on the word, the error detection code to be stored with the word.

(899) The control of block **3318** may be performed in response to an additional subsequent clock cycle of the CPU **102**, or the cache controller **220**.

(900) In response, the victim cache store queue **216** determines whether an additional write instruction is obtained. (Block **3322**). the event the victim cache store queue **216** determines another write instruction is obtained (e.g., the control of block **3322** returns a result of YES), the process returns to block **3302**. Alternatively, in the event the victim cache store queue **216** determines another write instruction is not obtained (e.g., the control of block **3322** returns a result of NO), the process **3300** may wait until a threshold timeout period occurs, thus ending the process **3300**.

(901) FIG. **34** is a block diagram of an example processor platform **3400** structured to execute the instructions of FIGS. **12-33** to implement the L1 cache **110** of FIGS. **1-5** and **10-11**. The processor platform **1000** can be, for example, a server, a personal computer, a workstation, a self-learning machine (e.g., a neural network), a mobile device (e.g., a cell phone, a smart phone, a tablet such as an iPad™), a personal digital assistant (PDA), an Internet appliance, a gaming console, or any other type of computing device.

(902) The processor platform **3400** of the illustrated example includes a processor **3412**. The

processor **3412** of the illustrated example is hardware. For example, the processor **3412** can be implemented by one or more integrated circuits, logic circuits, microprocessors, GPUs, DSPs, or controllers from any desired family or manufacturer. The hardware processor may be a semiconductor based (e.g., silicon based) device. In this example, the processor implements any element of the example L1 cache **110** as shown in FIGS. **1-5** and **10-11**.

(903) The processor **3412** of the illustrated example includes a local memory **3413** (e.g., a cache). The processor **3412** of the illustrated example is in communication with a main memory including a volatile memory **3414** and a non-volatile memory **3416** via a bus **3418**. The volatile memory **3414** may be implemented by Synchronous Dynamic Random Access Memory (SDRAM), Dynamic Random Access Memory (DRAM), RAMBUS® Dynamic Random Access Memory (RDRAM®) and/or any other type of random access memory device. The non-volatile memory **3416** may be implemented by flash memory and/or any other desired type of memory device. Access to the main memory **3414**, **3416** is controlled by a memory controller.

(904) The processor platform **3400** of the illustrated example also includes an interface circuit **3420**. The interface circuit **3420** may be implemented by any type of interface standard, such as an Ethernet interface, a universal serial bus (USB), a Bluetooth® interface, a near field communication (NFC) interface, and/or a PCI express interface.

(905) In the illustrated example, one or more input devices **3422** are connected to the interface circuit **3420**. The input device(s) **3422** permit(s) a user to enter data and/or commands into the processor **3412**. The input device(s) can be implemented by, for example, an audio sensor, a microphone, a camera (still or video), a keyboard, a button, a mouse, a touchscreen, a track-pad, a trackball, isopoint and/or a voice recognition system.

(906) One or more output devices **3424** are also connected to the interface circuit **3420** of the illustrated example. The output devices **3424** can be implemented, for example, by display devices (e.g., a light emitting diode (LED), an organic light emitting diode (OLED), a liquid crystal display (LCD), a cathode ray tube display (CRT), an in-place switching (IPS) display, a touchscreen, etc.), a tactile output device, a printer and/or speaker. The interface circuit **3420** of the illustrated example, thus, typically includes a graphics driver card, a graphics driver chip and/or a graphics driver processor.

(907) The interface circuit **3420** of the illustrated example also includes a communication device such as a transmitter, a receiver, a transceiver, a modem, a residential gateway, a wireless access point, and/or a network interface to facilitate exchange of data with external machines (e.g., computing devices of any kind) via a network **3426**. The communication can be via, for example, an Ethernet connection, a digital subscriber line (DSL) connection, a telephone line connection, a coaxial cable system, a satellite system, a line-of-site wireless system, a cellular telephone system, etc.

(908) The processor platform **3400** of the illustrated example also includes one or more mass storage devices **3428** for storing software and/or data. Examples of such mass storage devices **3428** include floppy disk drives, hard drive disks, compact disk drives, Blu-ray disk drives, redundant array of independent disks (RAID) systems, and digital versatile disk (DVD) drives.

(909) The machine executable instructions **3432** of FIGS. **12-33** may be stored in the mass storage device **3428**, in the volatile memory **3414**, in the non-volatile memory **3416**, and/or on a removable non-transitory computer readable storage medium such as a CD or DVD.

(910) From the foregoing, it will be appreciated that example methods, apparatus and articles of manufacture have been disclosed to facilitate write miss caching in cache system. The disclosed methods, apparatus and articles of manufacture improve the efficiency of using a computing device by reducing a data cache to reduce latency of a computing system and improving computer system efficiency to reduce stress on a computer core. The disclosed methods, apparatus and articles of manufacture are accordingly directed to one or more improvement(s) in the functioning of a computer.

(911) Although certain example methods, apparatus and articles of manufacture have been disclosed herein, the scope of coverage of this patent is not limited thereto. On the contrary, this patent covers all methods, apparatus and articles of manufacture fairly falling within the scope of the claims of this patent.

(912) The following claims are hereby incorporated into this Detailed Description by this reference, with each claim standing on its own as a separate embodiment of the present disclosure.

Claims

1. A device comprising: a cache memory; and a cache controller coupled to the cache memory and configured to: allocate a portion of the cache memory as a write miss cache; receive a write operation; determine whether the write operation is associated with a miss in a remainder of the cache memory; based on the write operation being associated with the miss in the remainder of the cache memory, cause a set of write miss information to be stored in the write miss cache; receive a read operation that is associated with the write operation; and service the read operation using the set of write miss information stored in the write miss cache.
2. The device of claim 1, wherein: the set of write miss information includes a set of data associated with the write operation; and the servicing of the read operation includes providing the set of data from the write miss cache.
3. The device of claim 1, wherein: the cache memory is a first cache memory; the device further comprises a second cache memory coupled to the cache controller and configured to be a main cache; and the cache controller is configured to: allocate the remainder of the cache memory as a victim cache; and cause the set of write miss information to be stored in the write miss cache further based on the write operation being associated with a miss in the main cache.
4. The device of claim 3, wherein: the main cache is a level-one (L1) main cache; the victim cache is an L1 victim cache; and the cache controller is an L1 cache controller.
5. The device of claim 4, wherein the cache controller is configured to: compare a utilization of the write miss cache to a threshold; and based on the utilization exceeding the threshold, cause the set of write miss information to be provides to a level-two (L2) cache subsystem.
6. The device of claim 5, wherein the threshold corresponds to a bandwidth of an interface of the L2 cache subsystem.
7. The device of claim 5, wherein the threshold corresponds to a size of the write miss cache.
8. The device of claim 1, wherein: the set of write miss information is a first set of write miss information; and the cache controller is configured to, when a first memory address of the first set of write miss information corresponds to a second memory address of a second set of write miss information stored in the write miss cache, to merge the first set of write miss information with the second set of write miss information.
9. The device of claim 1, wherein the write miss cache includes a byte enable register.
10. The device of claim 9, wherein the byte enable register is configured to store a value that specifies an element of the set of write miss information to be written.
11. A device comprising: a processing core configured to provide a write operation; a cache controller coupled to the processing core; a first cache memory coupled to the processing core; and a second cache memory coupled to the processing core that includes a write miss cache and a remainder; wherein the cache controller is configured to: determine whether the write operation is associated with a miss in the first cache memory and a miss in the remainder of the second cache memory; based on the write operation being associated with the miss in the first cache memory and the miss in the remainder of the second cache memory, cause a set of write miss information to be stored in the write miss cache; receive a read operation that is associated with the write operation; and service the read operation using the set of write miss information stored in the write miss cache.

12. A method comprising: allocating a portion of a cache memory as a write miss cache; receiving a write operation; determining whether the write operation is associated with a miss in a remainder of the cache memory; based on the write operation being associated with the miss in the remainder of the cache memory, storing a set of write miss information in the write miss cache; receiving a read operation that is associated with the write operation; and servicing the read operation using the set of write miss information stored in the write miss cache.
 13. The method of claim 12, wherein: the set of write miss information includes a set of data associated with the write operation; and servicing the read operation using the set of write miss information comprises providing the set of data from the write miss cache.
 14. The method of claim 12, wherein: the cache memory is a first cache memory; and the method further comprises: allocating the remainder of the first cache memory as a victim cache; and storing the set of write miss information in the write miss cache further based on the write operation being associated with a miss in a second cache memory.
 15. The method of claim 14, wherein: the second cache memory is a level-one (L1) main cache; and the first cache memory is an L1 victim cache.
 16. The method of claim 15, further comprising: comparing a utilization of the write miss cache to a threshold; and based on the utilization exceeding the threshold, providing the set of write miss information to a level-two (L2) cache subsystem.
 17. The method of claim 16, wherein the threshold corresponds to a bandwidth of an interface of the L2 cache subsystem.
 18. The method of claim 16, wherein the threshold corresponds to a size of the write miss cache.
 19. The method of claim 12, wherein: the set of write miss information is a first set of write miss information; and the method further comprises, when a first memory address of the first set of write miss information corresponds to a second memory address of a second set of write miss information stored in the write miss cache, merging the first set of write miss information with the second set of write miss information.
 20. The method of claim 12, wherein: the write miss cache includes a byte enable register; and the method further comprises storing a value in the byte enable register that specifies an element of the set of write miss information to be written.
-