

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12393563
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Lambert; Collin et al.

---

### Storage and structured search of historical security data

---

#### Abstract

A method includes ingesting event data over a network for a plurality of events obtained by disparate computing resources. Each event is associated with a respective timestamp and one or more ingestion-attributes. The method includes identifying whether the corresponding event is associated with any custom indexing-attributes defined by a user. The method also includes indexing the corresponding event into a data store as structured data based on the respective timestamp, the one or more ingestion-attributes, and any identified custom indexing-attributes. The method includes evicting any of the events of the event data in the data store for a period of time that satisfies an eviction time period threshold. The method also includes retrieving the data from the data store that is associated with the time range, the ingestion-attributes, or the one custom indexing-attributes.

---

**Inventors:** Lambert; Collin (Mountain View, CA), Basmov; Innokentiy (Mountain View, CA), Gaebel; Ethan Daniel (Mountain View, CA), Chang; Andrew Liang Ping (San Jose, CA), Ion; Iulia (Kirkland, WA)

**Applicant:** Google LLC (Mountain View, CA)

**Family ID:** 1000008765877

**Assignee:** Google LLC (Mountain View, CA)

**Appl. No.:** 17/451579

**Filed:** October 20, 2021

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20220035783 A1	Feb. 03, 2022

#### Related U.S. Application Data

## Publication Classification

**Int. Cl.:** **G06F7/00** (20060101); **G06F16/22** (20190101); **G06F16/27** (20190101); **G06F16/28** (20190101); **G06F16/951** (20190101); **G06F21/60** (20130101); **G06F21/62** (20130101)

**U.S. Cl.:**

**CPC** **G06F16/2228** (20190101); **G06F16/27** (20190101); **G06F16/284** (20190101); **G06F16/951** (20190101); **G06F21/604** (20130101); **G06F21/62** (20130101);

## Field of Classification Search

**CPC:** G06F (21/604); G06F (21/62); G06F (16/2477); G06F (16/1834); G06F (16/2228); G06F (16/27); G06F (16/284); G06F (16/951)

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6138121	12/1999	Costa et al.	N/A	N/A
8752178	12/2013	Coates et al.	N/A	N/A
10409817	12/2018	Dias	N/A	G06F 16/2477
10643002	12/2019	Veselov et al.	N/A	N/A
10776355	12/2019	Batsakis et al.	N/A	N/A
2003/0200192	12/2002	Bell	N/A	G06F 16/9537
2006/0075007	12/2005	Anderson et al.	N/A	N/A
2006/0265746	12/2005	Farley	726/13	G06F 21/577
2008/0208820	12/2007	Usey	N/A	G06F 16/313
2010/0114895	12/2009	Bhagwan et al.	N/A	N/A
2010/0185963	12/2009	Slik et al.	N/A	N/A
2012/0136921	12/2011	Samdadiya et al.	N/A	N/A
2015/0046251	12/2014	Smith	N/A	N/A
2015/0134795	12/2014	Theimer et al.	N/A	N/A
2015/0154249	12/2014	Dave	707/758	G06F 16/951
2016/0063001	12/2015	Scriffignano et al.	N/A	N/A
2016/0224600	12/2015	Munk	N/A	G06Q 20/08
2016/0299957	12/2015	A. C. et al.	N/A	N/A
2016/0306871	12/2015	Chauhan et al.	N/A	N/A
2016/0335361	12/2015	Teodorescu et al.	N/A	N/A
2017/0048261	12/2016	Gmach et al.	N/A	N/A
2017/0063920	12/2016	Thomas et al.	N/A	N/A
2017/0134243	12/2016	Levesque et al.	N/A	N/A
2018/0004826	12/2017	Reiner et al.	N/A	N/A
2018/0005274	12/2017	Calvillo et al.	N/A	N/A
2018/0069925	12/2017	Lavasani	N/A	N/A
2018/0089272	12/2017	Bath et al.	N/A	N/A

2018/0173583	12/2017	Braundmeier et al.	N/A	N/A
2018/0176244	12/2017	Gervais et al.	N/A	N/A
2018/0246926	12/2017	Altaf	N/A	G06F 16/215
2018/0276223	12/2017	Dhanasekaran	N/A	G06F 16/125
2018/0276232	12/2017	Dhanasekaran	N/A	G06F 16/25
2019/0018844	12/2018	Bhagwat et al.	N/A	N/A
2019/0018904	12/2018	Russell	N/A	G06F 16/27
2019/0073615	12/2018	Ronen et al.	N/A	N/A
2019/0095478	12/2018	Tankersley et al.	N/A	N/A
2019/0318019	12/2018	Abes et al.	N/A	N/A

## FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2000-040063	12/1999	JP	N/A
2017-522664	12/2016	JP	N/A
2018-523862	12/2017	JP	N/A
2015200051	12/2014	WO	N/A
2017031302	12/2016	WO	N/A

## OTHER PUBLICATIONS

International Search report for the related application No. PCT/US2019/053073 dated Sep. 26, 2019. cited by applicant

Japanese Office Action for the related Application No. 2022-135855, dated Sep. 28, 2023, 4 pages. cited by applicant

---

*Primary Examiner:* Uddin; Mohammed R

*Attorney, Agent or Firm:* Honigman LLP

---

## Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS (1) This U.S. patent application is a continuation of, and claims priority under 35 U.S.C. § 120 from, U.S. patent application Ser. No. 16/198,344, filed on Nov. 21, 2018. The disclosure of this prior art application is considered part of the disclosure of this application and is hereby incorporated by reference in its entirety.

### TECHNICAL FIELD

(1) This disclosure relates to a system for storage and structured search of historical security data.

### BACKGROUND

(2) Identity, access, data, and resource security incidents continue to grow rapidly. Similarly, the amount of raw security signals available to be collected and analyzed is also growing exponentially. These security signals are a critical resource that enterprises need to protect themselves from security breaches and downtime. This need is driving human analysts and security engineers to increase efficiency and prioritization when dealing with this resource.

### SUMMARY

(3) One aspect of the disclosure provides a method including ingesting, by data processing hardware, event data over a network for a plurality of events obtained by a plurality of disparate computing resources in communication with the data processing hardware. The event data includes a respective timestamp for each event of the event data that indicates a point in time when the event was obtained by one of the plurality of disparate computing resources. The event data also includes

at least one ingestion-attribute associated with each event of the event data, the at least one ingestion-attribute satisfying ingestion criteria required to permit ingesting of the associated event. For each of the plurality of events of the event data, the method includes identifying, by the data processing hardware, whether the corresponding event is associated with any custom indexing-attributes defined by a user for indexing events. The method also includes indexing, by the data processing hardware, the corresponding event into a data store as structured data based on the respective timestamp for the corresponding event, the at least one ingestion-attribute associated with the corresponding event, and any identified custom indexing-attributes associated with the corresponding event. The method also includes evicting, by the data processing hardware, any of the events of the event data that have been indexed into the data store as structured data for a period of time that satisfies an eviction time period threshold. The method also includes receiving, at the data processing hardware, a retrieval request for structured data stored in the data store, the retrieval request requesting structured data associated with at least one of a time range specified by the retrieval request, one or more ingestion-attributes specified by the retrieval request, or one or more custom indexing-attributes specified by the retrieval request. The method also includes retrieving, by the data processing hardware, the structured data from the data store that is associated with the at least one of the time range specified by the retrieval request, the one or more ingestion-attributes specified by the retrieval request, or the one or more custom indexing-attributes specified by the retrieval request.

(4) Implementations of the disclosure may include one or more of the following optional features. In some examples, the custom indexing-attributes defined by the user for indexing events each include a respective key-value pair defined by a customer of the plurality of disparate computing resources. In some examples, the method includes, for each of the plurality of events of the event data, applying, by the data processing hardware, a set of validity rules to determine whether the corresponding event is valid. When the corresponding event is valid based on the applied set of validity rules, the method includes indexing the corresponding event into the data store as structured data. When the corresponding event is invalid based on the applied set of validity rules, the method includes rejecting, by the data processing hardware, the corresponding event for indexing into the data store. The set of validity rules may include a set of priority rules to determine a priority of the corresponding event. In some implementations, when receiving the retrieval request, the method includes receiving a structured data retrieval offset, the structured data retrieval offset indicating a position in a list of structured data to be retrieved, and where only structured data after the position in the list of structured data is retrieved. The method may further including sending, by the data processing hardware, a portion of the retrieved structured data and a page token and the page token indicating a position in a list of the retrieved structured data. The portion of the retrieved structured data includes only data from earlier positions than the page token in the list. In some examples, the data store includes a distributed storage system. In other examples, the data store includes a relational database. At least one of the plurality of events of the event data may be indicative of a measured characteristic of a corresponding one of the plurality of disparate computing resources. A priority of the measured characteristic may be determined based on a set of priority rules. Optionally, the retrieval request requesting structured data is associated with a first time range and a second time range and the second time range different from the first time range. In some implementations, ingesting the event data includes obtaining the event data over the network from the plurality of disparate computing resources via an application programming interface. The method may further include receiving, at the data processing hardware, an eviction request for evicting data, and the eviction request for evicting data may be associated with at least one of a time range specified by the eviction request, one or more ingestion-attributes specified by the eviction request, or one or more custom indexing-attributes specified by the eviction request. The method may also include evicting, by the data processing hardware, the structured data from the data store that is associated with the at least one of a time range specified by the eviction request,

one or more ingestion-attributes specified by the eviction request, or one or more custom indexing-attributes specified by the eviction request. In some implementations, ingesting the event data is in response to at least one of: receiving an ingestion request, an indication from a time schedule, or an indication from an event. Retrieving the structured data may include verifying permissions of the structured data associated with the retrieval request.

(5) Another aspect of the disclosure provides a system including data processing hardware and memory hardware in communication with the data processing hardware. The memory hardware stores instruction that when executed on the data processing hardware cause the data processing hardware to perform operations. The operations include ingesting event data over a network for a plurality of events obtained by a plurality of disparate computing resources in communication with the data processing hardware. The event data includes a respective timestamp for each event of the event data that indicates a point in time when the event was obtained by one of the plurality of disparate computing resources. The event data also includes at least one ingestion-attribute associated with each event of the event data and the at least one ingestion-attribute satisfies ingestion criteria required to permit ingesting of the associated event. For each of the plurality of events of the event data, the operations include identifying whether the corresponding event is associated with any custom indexing-attributes defined by a user for indexing events. The operations also include indexing the corresponding event into a data store as structured data based on the respective timestamp for the corresponding event. The at least one ingestion-attribute is associated with the corresponding event and any identified custom indexing-attributes are also associated with the corresponding event. The operations also include evicting any of the events of the event data that have been indexed into the data store as structured data for a period of time that satisfies an eviction time period threshold. The operations also include receiving a retrieval request for structured data stored in the data store. The retrieval request requests structured data associated with at least one of a time range specified by the retrieval request, one or more ingestion-attributes specified by the retrieval request, or one or more custom indexing-attributes specified by the retrieval request. The operations also include retrieving the structured data from the data store that is associated with the at least one of the time range specified by the retrieval request, the one or more ingestion-attributes specified by the retrieval request, or the one or more custom indexing-attributes specified by the retrieval request.

(6) Implementations of the disclosure may include one or more of the following optional features. In some examples, the custom indexing-attributes defined by the user for indexing events each include a respective key-value pair defined by a customer of the plurality of disparate computing resources. The operations may further include, for each of the plurality of events of the event data, applying a set of validity rules to determine whether the corresponding event is valid. The operations may then also include that, when the corresponding event is valid based on the applied set of validity rules, indexing the corresponding event into the data store as structured data. The set of validity rules may include a set of priority rules to determine a priority of the corresponding event. In some implementations, the operations further include, when the corresponding event is invalid based on the applied set of validity rules, rejecting the corresponding event for indexing into the data store. The retrieval request may include receiving a structured data retrieval offset, and the structured data retrieval offset may indicate a position in a list of structured data to be retrieved. Only structured data after the position in the list of structured data may be retrieved. The operations, in some examples, further include sending a portion of the retrieved structured data. The operations then include sending a page token, and the page token indicates a position in a list of the retrieved structured data. The operations may also include where the portion of the retrieved structured data includes only data from earlier positions than the page token in the list. In some examples, the data store includes a distributed storage system. In other examples, the data store includes a relational database. The at least one of the plurality of events of the event data may be indicative of a measured characteristic of a corresponding one of the plurality of disparate

computing resources. The operations may include determining a priority of the measured characteristic based on a set of priority rules. The retrieval request requesting structured data, in some implementations, is associated with a first time range and a second time range, and the second time range different from the first time range. Ingesting the event data may include obtaining the event data over the network from the plurality of disparate computing resources via an application programming interface. In some implementations, the operations further include receiving an eviction request for evicting data, and the eviction request for evicting data is associated with at least one of a time range specified by the eviction request, one or more ingestion-attributes specified by the eviction request, or one or more custom indexing-attributes specified by the eviction request. The operations then include evicting the structured data from the data store that is associated with the at least one of a time range specified by the eviction request, one or more ingestion-attributes specified by the eviction request, or one or more custom indexing-attributes specified by the eviction request. In some implementations, ingesting the event data is in response to at least one of: receiving an ingestion request, an indication from a time schedule, or an indication from an event. Retrieving the structured data may include verifying permissions of the structured data associated with the retrieval request.

(7) The details of one or more implementations of the disclosure are set forth in the accompanying drawings and the description below. Other aspects, features, and advantages will be apparent from the description and drawings, and from the claims.

---

## Description

### DESCRIPTION OF DRAWINGS

(1) FIGS. 1A and 1B are schematic views of an example system for storing and searching structured event data.

(2) FIG. 2 is a schematic view of example components of an ingestion interface of the system of FIGS. 1A and 1B.

(3) FIG. 3 is a schematic view of example components of a persistence subsystem of the system of FIGS. 1A and 1B.

(4) FIG. 4 is a schematic view of example components of a retrieval interface of the systems of FIGS. 1A and 1B.

(5) FIG. 5 is a flowchart of an example method for storing and structuring event data.

(6) FIG. 6 is a schematic view of an example computing device that may be used to implement the systems and methods described herein.

(7) Like reference symbols in the various drawings indicate like elements.

### DETAILED DESCRIPTION

(8) As the amount of available raw security signals that must be collected and analyzed grows exponentially (e.g., security findings or events for resources across a vast distributed computing system), enterprises are searching for ways to increase efficiency in detecting and responding to security issues in a cloud environment.

(9) Implementations herein are directed toward systems and methods for enabling the indexing and joining of current and historic raw security data at-scale across previously disparate sources in-order to accelerate both programmatic and human analysis and their prioritization of the data to deliver insights and drive human understanding and prioritized response actions. In addition, implementations herein enable users to organize, manage, investigate, make informed decisions, and act based on abstractions of the user's assets, workloads and relevant threats, while further reducing the user's cognitive and effort load. These contributions reduce the time to detect and fix issues and minimizes the risks and blast radius of incidents by providing efficient ingestion and retrieval of structured time-stamped information about cloud resources and the security information

associated with the cloud resources.

(10) Referring to FIGS. 1A and 1B, in some implementations, an example system **100** includes a remote system **112**. The remote system **112** may be a single computer, multiple computers, or a distributed system (e.g., a cloud environment) having scalable/elastic computing resources **114** (e.g., data processing hardware **610** (FIG. 6)) and/or storage resources **116** (e.g., memory hardware **620** (FIG. 6)). The remote system **112** is connected to a plurality of disparate computing resources or clients **118**, **118a-n** through network **120**. A storage abstraction **150** (e.g., a distributed storage system or a data store) is overlain on the storage resources **116** to allow scalable use of the storage resources **116** by one or more of the client or computing resources **118**. The remote system **112** executes a structured data search system **160**. The search system **160** obtains and ingests event data **162** from the computing resources **118**. The event data **162** represents events associated with cloud resources. In some implementations, the event data **162** or associated event is indicative of a measured characteristic of a corresponding one of the plurality of disparate computing resources **118**. The event data **162**, for example, includes characteristics associated with cloud resources, security and privacy vulnerabilities, and caller-provided annotation data. In some implementations, the event data **162** forms a hierarchy with a parent-child relationship (e.g., a hierarchy of cloud resources). The system **160** may determine a priority of the measured characteristic based on a set of priority rules. The event data **162** may also form a graph with a one-to-many or many-to-many relationships between elements. For example, the event data **162** may represent a cloud resource and all security issues associated with the cloud resource.

(11) The event data **162** received by the search system **160** further includes a timestamp **164**, at least one ingestion-attribute **166**, and custom indexing-attributes **168**. The custom indexing-attributes **168** may or may not be included in the event data **162**. In some examples, the search system **160** includes an ingestion interface **200**, a persistence subsystem **300**, and a retrieval interface **400**. The ingestion interface **200** receives the event data **162**, processes the data **162**, and passes the ingested data **162** to the persistence subsystem **300**. The persistence subsystem **300** stores the event data **162** in the storage abstraction **150** as structured data. The storage abstraction **150** is configured to store the event data **162** from the computing resources **118**. The distributed storage system **150** may implement an archive **190** (e.g., tape archive) configured to back-up stored event data **162** for recovery purposes. The archive **190** may include a long retention period (e.g., three (3) years). The retrieval interface **400**, in response to a retrieval request **410** from a data requester **119**, delivers retrieval data **420** to the data requester **119**. The data requester **119** may be associated with a user/customer or entity that owns the event data **162** and corresponding retrieval data **420**, and therefore may access the retrieval data **420** to inspect the contents thereof by transmitting the retrieval request **410** specifying the contents to include in the retrieval data **420**.

(12) Referring to FIG. 1B, in some implementations, the distributed system **140** includes loosely coupled memory hosts **114**, **114a-n** (i.e., data processing hardware)), each having a computing resource **122** (e.g., one or more processors or central processing units (CPUs)) in communication with storage resources **116** (e.g., memory hardware, flash memory, dynamic random access memory (DRAM), phase change memory (PCM), and/or disks) that may be used for caching data. The storage abstraction **150** overlain on the storage resources **116** allows scalable use of the storage resources **116** by one or more clients **118**, **119**. The clients **118**, **119** may communicate with the memory hosts **114** through the network **120** (e.g., via remote procedure calls (RPC)). In some implementations, the remote distributed system **112** is “single-sided.” “Single-sided” refers to the method by which most of the request processing on the memory hosts **114** may be done in hardware rather than by software executed on CPUs **122** of the memory hosts **114**.

(13) The distributed system **112** may store event data **162** obtained from clients **118**, **119** into the storage resources **116** (e.g., storage abstraction **150**) of the remote memory hosts **114** and get the retrieval data **420** from the remote memory hosts **114** via network interface controllers (NIC) **126**. A network interface controller **126** (also known as a network interface card, network adapter, or

LAN adapter) may be a computer hardware component that connects a computing device/resource **122** to the network **120**. Both the memory hosts **114a-n** and the clients **118**, **119** may each have a network interface controller **126** for network communications. Each memory location **124** is configured to store event data **162**. As used herein, the clients **118**, **119** may include the disparate computing resources **118** that collect/obtain/measure the event data **162** ingested by the structured data system **160**, and the data requesters **119** associated with customers/users associated with the event data **162** ingested by the structured data system **160** and corresponding retrieval data **410** retrieved from the structured data system **160** in response to sending retrieval requests **410**.

(14) Referring now to FIG. 2, the ingestion interface **200** ingests the event data **162**. The ingestion interface **200** may ingest the event data **162** by receiving the event data **162** from the disparate computing resources **118** and/or by actively fetching the event data **162** over the network **120** (e.g., via an application programming interface (API)) from the disparate computing resources.

Optionally, the event data **162** may be written or updated via the API. That is, clients **118**, **119** may push the event data **162** to the ingestion interface **200** via the API. The ingestion interface **200** may ingest data in response to any number of stimuli. For example, ingesting the event data may be in response to receiving an ingestion request. That is, a user or client requests the ingestion interface **200** to ingest data. The ingestion interface **200** may also ingest data in response to an indication from a time schedule. For example, a time schedule for the ingestion interface may specify specific times or ranges of times that event data should be ingested. Optionally, the ingestion interface **200** may ingest data in response to an indication from an event. That is, ingesting specific event data may trigger the ingestion of additional event data **162** from the same or different client **118**.

(15) As previously discussed, the event data **162** includes the timestamp **164**. The timestamp **164** indicates a point in time when the respective computing resource **118** obtained the respective event. In some examples, the event is associated with a cloud resource and may include a characteristic or some other measurable parameter associated with operation of the cloud resource. For example, the event data **162** may include results from a security scanner that has scanned a remote computing resource **118**. In another example, the event data **162** includes a notification that a password has been set on a computing resource **118** associated with a remote database. The timestamp **164** may indicate a time when the event occurred. That is, referring back to the scanner example, the timestamp **164** may indicate when the security scanner completed the scan or the remote computing resource **118** received the scanning results from the scanner. In some implementations, the timestamp **164** may indicate a time when the event data **162** was ingested by the structured data search system **160**. In other implementations, the ingestion interface **200** provides a timestamp **164** (e.g., of the current time) when ingesting event data **162** that lacks a timestamp **164**.

(16) The event data **162** also includes one or more ingestion-attributes **166**. These attributes **166** are required for the ingestion interface **200** to permit ingesting of the data **162**. For example, the event data **162** may be compared against ingestion rules **210** (i.e., validity rules or ingestion criteria). If the rules **210** are not satisfied (e.g., missing an ingestion-attribute **166**), the data **162** may be determined invalid, and the data **162** may be discarded **212** (i.e., ignored or rejected) or otherwise not ingested. If the rules **210** are satisfied (e.g., all ingestion-attributes **166** are present and in proper format), the ingestion interface **200** may determine that the data **162** is valid ingest the data **162** and proceed in sending the ingested data **162** to the persistence subsystem **300** for indexing. In some implementations, the ingestion rules **210** verify more than just the presence of the ingestion-attributes **166**. For example, the ingestion rules **210** may enforce time and/or bandwidth constraints (e.g., an amount of data **162** or the rate at which the data **162** is obtained). In another example, the ingestion rules **210** enforce access control (i.e., permissions) to the data **162** and/or system **100** (i.e., verify that the data **162** has access or permission to the system **100** and/or that the system **100** has access or permission to the data **162**). In yet another example, the ingestion rules **210** enforce ordering and may reject out-of-order updates. The ingestion rules may include a set of priority rules that determine a priority of the corresponding event. The data may be ingested or indexed based in



part on the determined priority. The ingestion-attributes **166** may be strongly typed (i.e., the type and format of the attributes may be strictly enforced). The ingestion-attributes **166** may be represented as pairs (e.g., key, value).

(17) With continued reference to FIG. 2, the event data **162**, in some implementations, includes custom indexing-attributes **168**. The custom indexing-attributes **168** are not essential for the ingestion interface **200** to ingest the data **162**, but instead allow clients (e.g., data requesters **119**) further flexibility and customization when indexing and retrieving the event data **162**. That is, the event data **162** may be ingested regardless of the presence of custom indexing-attributes **168**. For example, a description attribute defined by the client may be optional and the ingestion interface **200** may still ingest event data **162** lacking the description. Like the ingestion-attributes **166**, the custom indexing-attributes **168** may be represented as (key, value) pairs that the customer or client define. In some implementations, the custom indexing-attributes **168** augment the event data **162** with additional caller-provided tuples (e.g., key, value, validity period).

(18) Referring now to FIG. 3, the ingested data **162** (still including timestamp **164**, at least one ingestion-attribute **166**, and any custom indexing-attributes **168**) is received by the persistence subsystem **300** as pre-indexed data **162a**. In some implementations, the persistence subsystem includes a data indexer **310**. The data indexer **310** accepts the pre-indexed data **162a** and identifies whether the pre-indexed data **162a** is associated with any custom-indexing-attributes **168** defined by the client **119** (or a user of the client). The data indexer **310** then indexes (i.e., structures) the ingested data **162** as structured data **162b** into the data store (e.g., storage abstraction) **150** based on the respective timestamp **164**, the ingestion-attributes **166**, and any custom indexing-attributes **168**. That is, the data indexer **310** orders and organizes the data **162** for efficient updating and retrieval in persistent storage. In some implementations, the data store **150** is associated with a relational database and the data indexer **310** indexes the data **162** into the relational database. Additionally or alternatively, the data store **150** may be associated with a distributed database and the data indexer **310** indexes the data **162b** into the distributed database.

(19) In some implementations, the data indexer **310**, when structuring the data **162** (e.g., pre-indexed data **162a**), adds the custom indexing-attributes **168**. For example, the data indexer **310** may add a timeline attribute that includes all changes in the respective data's **162** attributes, properties, augmentations, and relationship structure (e.g., parent-child relationship), etc. over a specified period of time.

(20) In some examples, the persistence subsystem **300** includes a data evictor **320**. The data evictor **320** may include retention thresholds **322**. The persistence subsystem **300**, through the data evictor **320**, may limit the amount of time during which structured data **162b** is stored in the data store **150**. That is, structured data **162b** may have a respective retention threshold or eviction time period thresholds **322** (e.g., three months) and when an amount of time the data **162b** has been stored in data store **150** satisfies this threshold **322**, the data **162b** may be evicted (i.e., deleted) from the data store **150**. Different elements of the data **162** may have different retention thresholds **322**. The threshold that is applied to an element of structured data **162b** may be dependent upon one or more of: the respective timestamp **164**, the ingestion-attributes **166**, and any custom indexing-attributes **168**. For example, an ingestion-attribute **166** may include an owner of the structured data **162b** (i.e., the owner of the cloud resource associated with the event of the structured data **162b**). The retention threshold **322** may be assigned to structured data **162b** based on the status of the owner ingestion-attribute **166**.

(21) In some implementations, structured data **162b** has more than one respective retention threshold **322**. For example, structured data **162b** may be associated with a last modified retention threshold **324** and a total retention threshold **326**. Structured data **162b** may, at times, be an update to previously received ingested data **162**. For instance, structured data **162b** may update the status (e.g., availability) of a cloud resource. The subsystem **300** may store older “versions” of the structured data **162b** for less time than the total retention threshold **326** (i.e., less time than the

“latest” version). For example, the last modified retention threshold **324** may be three months while the total retention threshold **326** may be thirteen months. When structured data **162b** is updated, the older version of the data **162** (i.e., the data before it was updated) may be maintained for the last modified retention threshold **324** (three months) while the latest update **162b** may be maintained for the total retention threshold **326** (thirteen months) unless later updated again.

(22) With continued reference to FIG. **3**, in some implementations, the persistence subsystem **300** receives an eviction request **328** from a user or client of the structured data search system **160**. For instance, the user or client sending the eviction request **328** may be the same user or client associated with the data requester **119**. The eviction request **328** requests to evict structured data **162b** from the data store **150** that is associated with at least one of: a time range **165** specified by the eviction request **328**, one or more ingestion-attributes **166** specified by the eviction request **328**, or one or more custom indexing-attributes **168** specified by the eviction request **328**. The data evictor **320** may evict any amount of structured data **162b** based on the eviction request **328**. For example, a user or client **119** may request eviction of all data associated with the respective client **119** as specified by a corresponding ingestion-attribute **166**. In another example, a client **119** may request eviction of all data having timestamps **164** that fall within the time range **165** specified by the eviction request **328**. The data evictor **320** may require verification (e.g., a username and password) that the requester **119** is authorized to evict the requested data **162**.

(23) Referring now to FIG. **4**, the retrieval interface **400** (i.e., a query interface) receives a retrieval request **410** for structured data **162b** stored in the data store **150**. The retrieval request **410** may be substantially similar to the eviction request **328** of FIG. **3** except that the retrieval request **410** is requesting retrieval of the structured data **162b** as retrieval data **420** from the data store **150**. Thus, the retrieval request **410** may request structured data **162b** from the data store **150** that is associated with at least one of: the time range **165** specified by the retrieval request **410**, one or more ingestion-attributes **166** specified by the retrieval request **410**, or one or more custom indexing-attributes **168** specified by the retrieval request **410**. The custom indexing-attributes **168** specified by the retrieval request **410** may include attributes generated by the data indexer **310** (e.g., the timeline attribute). In some examples, the retrieval request **410** includes a request retrieve data from a specific point in time or to compare data **162b** between a first time range and a second time range (also known as a “diff”), where each time range is different. The diff may be between more than two time ranges. The comparison may produce additional attributes. For example, the comparison may produce an attribute indicating whether data **162b** existed at or both points in time or whether any of the data's **162b** attributes, properties, or augmentations have changed between two points in time. In another example, the retrieval request **410** may request all data **162b** associated with a specific cloud resource **118**. The retrieval interface **400** will fetch the requested structured data **162b** from the data store **150** and return the data **162b** as retrieval data **420** to the requester **119**. In some implementations, the retrieval interface **400** sorts and/or groups the retrieval data **420** based on attributes **166**, **168** and/or timestamps **164**.

(24) The retrieval interface **400** may also filter, sort, or group the retrieval data **420** on any elements of the data **420**, such as augmentations and relationship structure (e.g., parent-child relationship). The filtering may be specified in a standardized or proprietary query language. The retrieval interface may further group the retrieval data **420** with an aggregation function. For example, the aggregation function may include count, sum, and/or average. In some implementations, retrieval requests **410** may be “chained” or otherwise sequenced together to expose clustering, correlation, and causality between structured event data **162b** associated with respective timestamps **164**, attributes **166**, **168**, or any other events.

(25) In some implementations, receiving the retrieval request **410** includes receiving a structured data retrieval offset **412**. The structured data retrieval offset **412** indicates a position in a list of structured data **162b** to be retrieved. Only structured data **162b** after the position in the list of structured data **162b** is then retrieved. For example, if the retrieval request **410** includes a

structured data retrieval offset **412** of fifty (50), and the retrieval interface **400** fetches a list of one-hundred (100) elements of structured data **162b** that correspond to the retrieval request **410**, the retrieval interface **400** returns elements fifty (50) through one-hundred (100) instead of all of the data elements. In other implementations, the structured data retrieval offset **412** indicates a quantity of elements of the structured data **162b** to return at a time. For example, the retrieval request **410** may request that the retrieval interface **400** only returns ten (10) results at a time. Optionally, the retrieval data **420** returned by the retrieval interface **400** may include a page token **422**. When the retrieval interface **400** returns only a portion of the retrieval data **420** that corresponds with the retrieval request **410**, the page token **422** may indicate a position in a list that corresponds with the data **420** that has been returned. That is, if the retrieval interface **400** has one-hundred (100) elements of retrieval data **420** in a list to return, and only returns ten (10) elements, the page token **422** may indicate that the next element of the retrieval data **420** that the retrieval interface **400** will return is the eleventh element. A follow-up retrieval request **410** may then include the page token **422** to indicate to the retrieval interface **400** that the requester is ready for the next “batch” or group of data **420**.

(26) Thus, the system **100** enables customers and clients to secure themselves at a high level of abstraction. That is, assets, vulnerabilities, threat and risk assessments and detections are prioritized and personalized at scale to the relevant class of business and contextualized applications (or workloads, services, etc.) that the respective customer or client has deployed. Instead of tactically securing individual resources, the system **100** secures the client's environment holistically. Specifically, the system **100** enables both automatic and assisted discovery of declared and inferred relationships for the workload and its underlying services and resources at scale. The system **100** enables both automatic and assisted understanding and baselining of the static and dynamic behavior and relationship changes normal to the specific workload (or workload class) at scale. The system **100** also enables both automatic and assisted detection of stat and dynamic anomalies at scale and the understanding of types and/or values of target data present in the context of specific applications, workloads, and workload classes. The system **100** permits either automatic or assisted targeting and findings prioritization mapped to common threat actors methods and vulnerabilities personalized for a specific client or customer, the client's specific business context, applications, workloads, and workload classes.

(27) FIG. 5 is a flowchart of an example method **500** for storing and structuring event data. The flowchart starts at operation **502** by ingesting, by data processing hardware **114**, event data **162** over a network **120** for a plurality of events obtained by a plurality of disparate computing resources **118** in communication with the data processing hardware **114**. In some implementations, ingesting the event data is in response to at least one of: receiving an ingestion request, an indication from a time schedule, or an indication from an event. The event data **162** includes a respective timestamp **164** for each event of the event data **162** that indicates a point in time when the event was obtained by one of the plurality of disparate computing resources **118**. The event data **162** also includes at least one ingestion-attribute **166** associated with each event of the event data **162**. The at least one ingestion-attribute **166** satisfies ingestion criteria **210** required to permit ingesting of the associated event. In some implementations, ingesting the event data **162** includes fetching the event data **162** over the network **120** from the plurality of disparate computing resources **118** via an application programming interface.

(28) For each of the plurality of events of the event data **162**, the method **500**, at step **504**, includes identifying, by the data processing hardware **114**, whether the corresponding event is associated with any custom indexing-attributes **168** defined by a user for indexing events. In some examples, the custom indexing-attributes **168** are defined by the user for indexing events that each include a respective key-value pair defined by a customer of the plurality of disparate computing resources **118**. At step **506**, the method **500** also includes for each of the plurality of events of the event data **162**, indexing, by the data processing hardware **114**, the corresponding event into a data store **150**

as structured data **162b** based on the respective timestamp **164** for the corresponding event, the at least one ingestion-attribute **166** associated with the corresponding event, and any identified custom indexing-attributes **168** associated with the corresponding event. In some examples, the data store **150** includes a distributed storage system. In other examples, the data store **150** includes a relational database.

(29) At step **508**, the method **500** includes evicting, by the data processing hardware **114**, any of the events of the event data **162** that have been indexed into the data store **150** as structured data **162b** for a period of time that satisfies an eviction time period threshold **322**. The method **500**, at step **510**, includes receiving, at the data processing hardware **114**, a retrieval request **410** for structured data **162b** stored in the data store **150**, the retrieval request **410** requesting structured data **162b** associated with at least one of a time range **165** specified by the retrieval request **410**, one or more ingestion-attributes **166** specified by the retrieval request **410**, or one or more custom indexing-attributes **168** specified by the retrieval request **410**. The retrieval request **410** requesting structured data **162b** may be associated with a first time range and a second time range. The second time range is different from the first time range.

(30) At step **512**, the method **500** includes retrieving, by the data processing hardware **114**, the structured data **162b** from the data store **150** that is associated with the at least one of the time range **165** specified by the retrieval request **410**, the one or more ingestion-attributes **166** specified by the retrieval request **410**, or the one or more custom indexing-attributes **168** specified by the retrieval request **410**. For instance, the structured data **162b** may be associated with the time range **165** specified by the retrieval request **410** when the structured data **162b** includes a corresponding timestamp **164** that falls within the specified time range **165**. In some implementations, receiving the retrieval request **410** includes receiving a structured data retrieval offset **412**, the structured data retrieval offset **412** indicating a position in a list of structured data **162b** to be retrieved, and wherein only structured data **162b** after the position in the list of structured data **162b** is retrieved. Retrieving the structured data **162b** may include verifying permissions of the structured data associated with the retrieval request **410**.

(31) Optionally, the method **500** includes for each of the plurality of events of the event data **162**, applying, by the data processing hardware **114**, a set of validity rules **210** to determine whether the corresponding event is valid. The method **500** may also include, when the corresponding event is valid based on the applied set of validity rules **210**, indexing the corresponding event into the data store **150** as structured data **162b**. The set of validity rules **210** may include a set of priority rules to determine a priority of the corresponding event. When the corresponding event is invalid based on the applied set of validity rules, the method **500** may include rejecting, by the data processing hardware **114**, the corresponding event for indexing into the data store **150**. In some examples, the method **500** includes sending, by the data processing hardware **114**, a portion of the retrieved structured data **162b** and a page token **422**. The page token **422** indicates a position in a list of the retrieved structured data **162b** and the portion of the retrieved structured data **162b** includes only data **162b** from earlier positions than the page token **422** in the list. At least one of the plurality of events of the event data **162** may be indicative of a measured characteristic of a corresponding one of the plurality of disparate computing resources **118**. Optionally, the method **500** includes determining a priority of the measured characteristic based on a set of priority rules.

(32) A software application (i.e., a software resource) may refer to computer software that causes a computing device to perform a task. In some examples, a software application may be referred to as an “application,” an “app,” or a “program.” Example applications include, but are not limited to, system diagnostic applications, system management applications, system maintenance applications, word processing applications, spreadsheet applications, messaging applications, media streaming applications, social networking applications, and gaming applications.

(33) FIG. **6** is schematic view of an example computing device **600** that may be used to implement the systems and methods described in this document. The computing device **600** is intended to

represent various forms of digital computers, such as laptops, desktops, workstations, personal digital assistants, servers, blade servers, mainframes, and other appropriate computers. The components shown here, their connections and relationships, and their functions, are meant to be exemplary only, and are not meant to limit implementations of the inventions described and/or claimed in this document.

(34) The computing device **600** includes a processor **610**, memory **620**, a storage device **630**, a high-speed interface/controller **640** connecting to the memory **620** and high-speed expansion ports **650**, and a low speed interface/controller **660** connecting to a low speed bus **670** and a storage device **630**. Each of the components **610**, **620**, **630**, **640**, **650**, and **660**, are interconnected using various busses, and may be mounted on a common motherboard or in other manners as appropriate. The processor **610** can process instructions for execution within the computing device **600**, including instructions stored in the memory **620** or on the storage device **630** to display graphical information for a graphical user interface (GUI) on an external input/output device, such as display **680** coupled to high speed interface **640**. In other implementations, multiple processors and/or multiple buses may be used, as appropriate, along with multiple memories and types of memory. Also, multiple computing devices **600** may be connected, with each device providing portions of the necessary operations (e.g., as a server bank, a group of blade servers, or a multi-processor system).

(35) The memory **620** stores information non-transitorily within the computing device **600**. The memory **620** may be a computer-readable medium, a volatile memory unit(s), or non-volatile memory unit(s). The non-transitory memory **620** may be physical devices used to store programs (e.g., sequences of instructions) or data (e.g., program state information) on a temporary or permanent basis for use by the computing device **600**. Examples of non-volatile memory include, but are not limited to, flash memory and read-only memory (ROM)/programmable read-only memory (PROM)/erasable programmable read-only memory (EPROM)/electronically erasable programmable read-only memory (EEPROM) (e.g., typically used for firmware, such as boot programs). Examples of volatile memory include, but are not limited to, random access memory (RAM), dynamic random access memory (DRAM), static random access memory (SRAM), phase change memory (PCM) as well as disks or tapes.

(36) The storage device **630** is capable of providing mass storage for the computing device **600**. In some implementations, the storage device **630** is a computer-readable medium. In various different implementations, the storage device **630** may be a floppy disk device, a hard disk device, an optical disk device, or a tape device, a flash memory or other similar solid state memory device, or an array of devices, including devices in a storage area network or other configurations. In additional implementations, a computer program product is tangibly embodied in an information carrier. The computer program product contains instructions that, when executed, perform one or more methods, such as those described above. The information carrier is a computer- or machine-readable medium, such as the memory **620**, the storage device **630**, or memory on processor **610**.

(37) The high speed controller **640** manages bandwidth-intensive operations for the computing device **600**, while the low speed controller **660** manages lower bandwidth-intensive operations. Such allocation of duties is exemplary only. In some implementations, the high-speed controller **640** is coupled to the memory **620**, the display **680** (e.g., through a graphics processor or accelerator), and to the high-speed expansion ports **650**, which may accept various expansion cards (not shown). In some implementations, the low-speed controller **660** is coupled to the storage device **630** and a low-speed expansion port **690**. The low-speed expansion port **690**, which may include various communication ports (e.g., USB, Bluetooth, Ethernet, wireless Ethernet), may be coupled to one or more input/output devices, such as a keyboard, a pointing device, a scanner, or a networking device such as a switch or router, e.g., through a network adapter.

(38) The computing device **600** may be implemented in a number of different forms, as shown in the figure. For example, it may be implemented as a standard server **600a** or multiple times in a

group of such servers **600a**, as a laptop computer **600b**, or as part of a rack server system **600c**.

(39) Various implementations of the systems and techniques described herein can be realized in digital electronic and/or optical circuitry, integrated circuitry, specially designed ASICs (application specific integrated circuits), computer hardware, firmware, software, and/or combinations thereof. These various implementations can include implementation in one or more computer programs that are executable and/or interpretable on a programmable system including at least one programmable processor, which may be special or general purpose, coupled to receive data and instructions from, and to transmit data and instructions to, a storage system, at least one input device, and at least one output device.

(40) These computer programs (also known as programs, software, software applications or code) include machine instructions for a programmable processor, and can be implemented in a high-level procedural and/or object-oriented programming language, and/or in assembly/machine language. As used herein, the terms “machine-readable medium” and “computer-readable medium” refer to any computer program product, non-transitory computer readable medium, apparatus and/or device (e.g., magnetic discs, optical disks, memory, Programmable Logic Devices (PLDs)) used to provide machine instructions and/or data to a programmable processor, including a machine-readable medium that receives machine instructions as a machine-readable signal. The term “machine-readable signal” refers to any signal used to provide machine instructions and/or data to a programmable processor.

(41) The processes and logic flows described in this specification can be performed by one or more programmable processors, also referred to as data processing hardware, executing one or more computer programs to perform functions by operating on input data and generating output. The processes and logic flows can also be performed by special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit). Processors suitable for the execution of a computer program include, by way of example, both general and special purpose microprocessors, and any one or more processors of any kind of digital computer. Generally, a processor will receive instructions and data from a read only memory or a random access memory or both. The essential elements of a computer are a processor for performing instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto optical disks, or optical disks. However, a computer need not have such devices. Computer readable media suitable for storing computer program instructions and data include all forms of non-volatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto optical disks; and CD ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

(42) To provide for interaction with a user, one or more aspects of the disclosure can be implemented on a computer having a display device, e.g., a CRT (cathode ray tube), LCD (liquid crystal display) monitor, or touch screen for displaying information to the user and optionally a keyboard and a pointing device, e.g., a mouse or a trackball, by which the user can provide input to the computer. Other kinds of devices can be used to provide interaction with a user as well; for example, feedback provided to the user can be any form of sensory feedback, e.g., visual feedback, auditory feedback, or tactile feedback; and input from the user can be received in any form, including acoustic, speech, or tactile input. In addition, a computer can interact with a user by sending documents to and receiving documents from a device that is used by the user; for example, by sending web pages to a web browser on a user's client device in response to requests received from the web browser.

(43) A number of implementations have been described. Nevertheless, it will be understood that

various modifications may be made without departing from the spirit and scope of the disclosure. Accordingly, other implementations are within the scope of the following claims.

## Claims

1. A computer-implemented method executed on data processing hardware that causes the data processing hardware to perform operations comprising: receiving an event over a network obtained by a plurality of disparate computing resources in communication with the data processing hardware, the event comprising a corresponding ingestion-attribute associated with the event; determining that the corresponding ingestion-attribute associated with the event satisfies ingestion criteria; and in response to determining that the corresponding ingestion-attribute associated with the event satisfies the ingestion criteria: ingesting the event obtained by the one of the plurality of disparate computing resources for persistent storage in a data store by indexing, using the corresponding ingestion-attribute as an indexing-attribute, the event as structured data at the data store, the data store associated with an eviction time period threshold representing a maximum amount of time that an event may be stored at the data store and at least one attribute for determining whether the eviction time period threshold applies to a particular event stored in the data store based on a corresponding ingestion-attribute associated with the particular event; determining, based on the corresponding ingestion-attribute matching the at least one attribute, that the eviction time period threshold applies to the event; and based on determining that the eviction time period threshold applies to the event, applying the eviction time period threshold to the event by: determining that the event has been stored in the data store for a period of time that satisfies the eviction time period threshold; and in response to determining that the event has been stored in the data store for the period of time that satisfies the eviction time period threshold, evicting the event from the data store.
2. The computer-implemented method of claim 1, wherein ingesting the event is in response to at least one of: receiving an ingestion request; an indication from a time schedule; or an indication from an event.
3. The computer-implemented method of claim 1, wherein the operations further comprise, identifying whether the event is associated with any custom indexing-attributes defined by a user for indexing events.
4. The computer-implemented method of claim 1, wherein the operations further comprise: applying a set of validity rules to the event; determining whether the event is valid based on the set of validity rules; and when the event is valid, indexing the event into the data store as structured data.
5. The computer-implemented method of claim 4, wherein the set of validity rules comprises a set of priority rules to determine a priority of the event.
6. The computer-implemented method of claim 4, wherein the operations further comprise, when the event is invalid, rejecting the event for indexing into the data store.
7. The computer-implemented method of claim 1, wherein the operations further comprise: receiving a retrieval request for data stored in the data store, wherein receiving the retrieval request comprises receiving a data retrieval offset, the data retrieval offset indicating a position in a list of data to be retrieved, and wherein only data after the position in the list of data is retrieved.
8. The computer-implemented method of claim 1, wherein the data store comprises a distributed storage system.
9. The computer-implemented method of claim 1, wherein ingesting the event comprises obtaining the event by the one of the plurality of disparate computing resources via an application programming interface.
10. The computer-implemented method of claim 1, wherein the event is indicative of a measured characteristic of a corresponding one of the plurality of disparate computing resources.

11. A system comprising: data processing hardware; and memory hardware in communication with the data processing hardware, the memory hardware storing instructions that, when executed on the data processing hardware, cause the data processing hardware to perform operations comprising: receiving an event over a network obtained by a plurality of disparate computing resources in communication with the data processing hardware, the event comprising a corresponding ingestion-attribute associated with the event; determining that the corresponding ingestion-attribute associated with the event satisfies ingestion criteria; and in response to determining that the corresponding ingestion-attribute associated with the event satisfies the ingestion criteria: ingesting the event obtained by the one of the plurality of disparate computing resources for persistent storage in a data store by indexing, using the corresponding ingestion-attribute as an indexing-attribute, the event as structured data at the data store, the data store associated with an eviction time period threshold representing a maximum amount of time that an event may be stored at the data store and at least one attribute for determining whether the eviction time period threshold applies to a particular event stored in the data store based on a corresponding ingestion-attribute associated with the particular event; determining, based on the corresponding ingestion-attribute matching the at least one attribute, that the eviction time period threshold applies to the event; and based on determining that the eviction time period threshold applies to the event, applying the eviction time period threshold to the event by: determining that the event has been stored in the data store for a period of time that satisfies the eviction time period threshold; and in response to determining that the event has been stored in the data store for the period of time that satisfies the eviction time period threshold, evicting the event from the data store.
12. The system of claim 11, wherein ingesting the event is in response to at least one of: receiving an ingestion request; an indication from a time schedule; or an indication from an event.
13. The system of claim 11, wherein the operations further comprise, identifying whether the event is associated with any custom indexing-attributes defined by a user for indexing events.
14. The system of claim 11, wherein the operations further comprise: applying a set of validity rules to the event; determining whether the event is valid based on the set of validity rules; and when the event is valid, indexing the event into the data store as structured data.
15. The system of claim 14, wherein the set of validity rules comprises a set of priority rules to determine a priority of the event.
16. The system of claim 14, wherein the operations further comprise, when the event is invalid, rejecting the event for indexing into the data store.
17. The system of claim 11, wherein the operations further comprise: receiving a retrieval request for data stored in the data store, wherein receiving the retrieval request comprises receiving a data retrieval offset, the data retrieval offset indicating a position in a list of data to be retrieved, and wherein only data after the position in the list of data is retrieved.
18. The system of claim 11, wherein the data store comprises a distributed storage system.
19. The system of claim 11, wherein ingesting the event comprises obtaining the event by the one of the plurality of disparate computing resources via an application programming interface.
20. The system of claim 11, wherein the event is indicative of a measured characteristic of a corresponding one of the plurality of disparate computing resources.
-