



US012393567B2

(12) **United States Patent**
Shuster

(10) **Patent No.:** **US 12,393,567 B2**
(45) **Date of Patent:** ***Aug. 19, 2025**

(54) **LOGFILE COLLECTION AND
CONSOLIDATION**

(56) **References Cited**

U.S. PATENT DOCUMENTS

(71) Applicant: **Red Hat, Inc.**, Raleigh, NC (US)
(72) Inventor: **Boaz Shuster**, Kfar Saba (IL)
(73) Assignee: **Red Hat, Inc.**, Raleigh, NC (US)
(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 45 days.
This patent is subject to a terminal disclaimer.
(21) Appl. No.: **17/532,023**
(22) Filed: **Nov. 22, 2021**

5,978,940	A	11/1999	Newman et al.
5,987,633	A	11/1999	Newman et al.
6,374,296	B1	4/2002	Lim et al.
6,789,156	B1	9/2004	Waldspurger
7,209,898	B2	4/2007	Pfeiffer et al.
7,653,836	B1	1/2010	Chatterjee et al.
8,464,335	B1	6/2013	Sinha et al.
9,736,243	B2	8/2017	Chablani et al.
9,959,489	B2	5/2018	Yamamoto et al.
10,142,195	B1	11/2018	Sledz et al.
10,181,991	B1	1/2019	Mathews et al.
10,387,370	B2	8/2019	Shuster
2003/0009486	A1	1/2003	Shimada et al.
2003/0105811	A1	6/2003	Laborde et al.
2003/0135507	A1	7/2003	Hind et al.

(Continued)

(65) **Prior Publication Data**

US 2022/0075769 A1 Mar. 10, 2022

Related U.S. Application Data

(63) Continuation of application No. 15/202,986, filed on Jul. 6, 2016, now Pat. No. 11,188,513.

(51) **Int. Cl.**

G06F 16/23 (2019.01)
G06F 16/18 (2019.01)
G06F 16/182 (2019.01)
H04L 43/0817 (2022.01)

(52) **U.S. Cl.**

CPC **G06F 16/2322** (2019.01); **G06F 16/1805** (2019.01); **G06F 16/2358** (2019.01); **G06F 16/182** (2019.01); **H04L 43/0817** (2013.01)

(58) **Field of Classification Search**

CPC G06F 16/2322; G06F 16/2358; G06F 16/1805; G06F 16/182

See application file for complete search history.

OTHER PUBLICATIONS

Non-Final Office Action for U.S. Appl. No. 15/202,986, mailed Jun. 1, 2019, 23 pages.

(Continued)

Primary Examiner — Amresh Singh

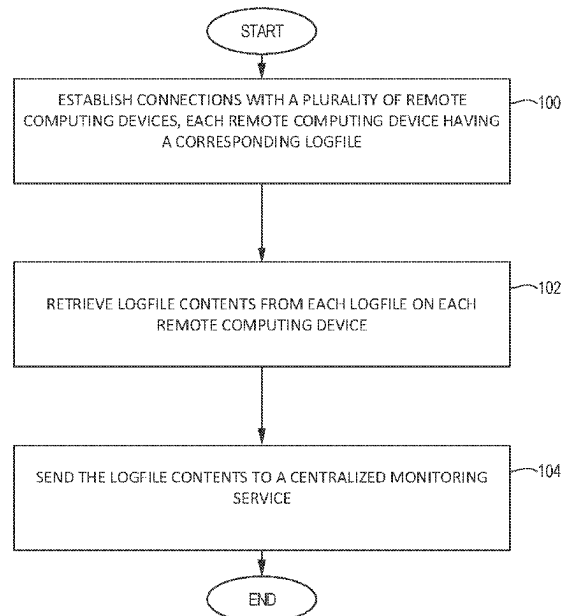
(74) *Attorney, Agent, or Firm* — Dority & Manning, P.A.

(57)

ABSTRACT

Mechanisms for consolidating log information from remote computing devices are provided. Connections with a plurality of remote computing devices are established. Each remote computing device has a corresponding logfile. For a plurality of iterations, logfile contents from each logfile on each remote computing device are retrieved, and the logfile contents are sent to a centralized monitoring service.

13 Claims, 6 Drawing Sheets



(56)

References Cited

U.S. PATENT DOCUMENTS

2004/0078722	A1	4/2004	Pfeiffer et al.	
2005/0235058	A1	10/2005	Rackus et al.	
2005/0273667	A1	12/2005	Shrivastava et al.	
2006/0106851	A1 *	5/2006	Warshawsky	G06F 16/2358
2007/0050777	A1	3/2007	Hutchinson et al.	
2007/0100991	A1	5/2007	Daniels et al.	
2007/0283194	A1 *	12/2007	Villella	G06F 11/3476 714/E11.204
2009/0013149	A1	1/2009	Uhlig et al.	
2009/0138703	A1	5/2009	Schneider	
2009/0313503	A1	12/2009	Atluri et al.	
2011/0153616	A1	6/2011	Torres	
2011/0208663	A1 *	8/2011	Kennis	G06Q 30/018 705/317
2011/0276695	A1	11/2011	Maldaner	
2012/0017028	A1	1/2012	Tsirkin	
2012/0216007	A1 *	8/2012	Tsirkin	G06F 9/526 711/E12.002
2013/0227352	A1	8/2013	Kumarasamy et al.	
2013/0332424	A1 *	12/2013	Nos	G06F 16/1734 707/E17.014
2014/0040180	A1	2/2014	Ruhl et al.	
2014/0058801	A1	2/2014	Deodhar et al.	
2014/0181036	A1 *	6/2014	Dhamankar	G06F 16/1734 707/648
2014/0181461	A1	6/2014	Kegel et al.	
2014/0279920	A1 *	9/2014	Madhavarapu	G06F 16/2358 707/649
2014/0282031	A1	9/2014	Hinterbichler et al.	
2015/0220605	A1	8/2015	Syed et al.	
2015/0256423	A1	9/2015	Stearns	
2015/0269183	A1 *	9/2015	Avati	G06F 16/1865 707/634
2015/0269213	A1 *	9/2015	Avati	G06F 16/9017 707/693
2015/0286663	A1	10/2015	Mangan, III et al.	
2015/0324297	A1	11/2015	Tati et al.	
2016/0041976	A1 *	2/2016	Pradeep	G06F 16/2358 707/695
2016/0275101	A1 *	9/2016	Hinterbichler	G06F 16/1734
2016/0275128	A1 *	9/2016	Dayka	G06F 16/2358
2016/0292263	A1	10/2016	Ferrar	
2016/0299934	A1 *	10/2016	Karandikar	G06F 16/248
2016/0378577	A1	12/2016	Russell et al.	
2017/0075922	A1 *	3/2017	Torman	G06F 16/1805
2017/0104820	A1	4/2017	Golander et al.	
2017/0331880	A1 *	11/2017	Crofton	G06F 16/182
2018/0013612	A1	1/2018	Shuster	

OTHER PUBLICATIONS

Final Office Action for U.S. Appl. No. 15/202,986, mailed Jan. 18, 2019, 18 pages.

Advisory Action for U.S. Appl. No. 15/202,986, mailed May 23, 2019, 5 pages.

Non-Final Office Action for U.S. Appl. No. 15/202,986, mailed Sep. 17, 2019, 20 pages.

Final Office Action for U.S. Appl. No. 15/202,986, mailed Apr. 22, 2020, 25 pages.

Advisory Action for U.S. Appl. No. 15/202,986, mailed Jul. 17, 2020, 3 pages.

Non-Final Office Action for U.S. Appl. No. 15/202,986, mailed Jan. 14, 2021, 26 pages.

Notice of Allowance for U.S. Appl. No. 15/202,986, mailed Jul. 30, 2021, 13 pages.

* cited by examiner


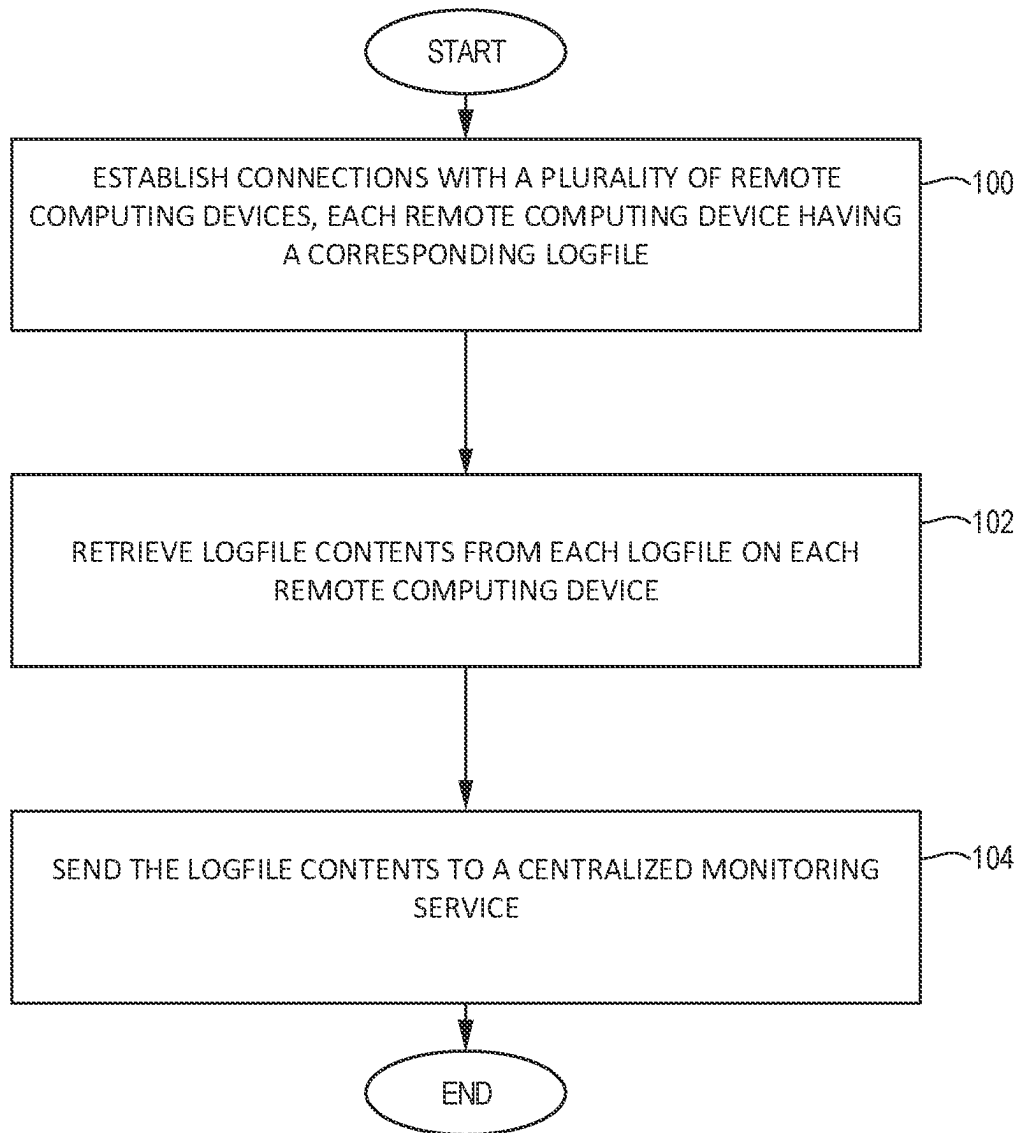



FIG. 1

*FIG. 2*



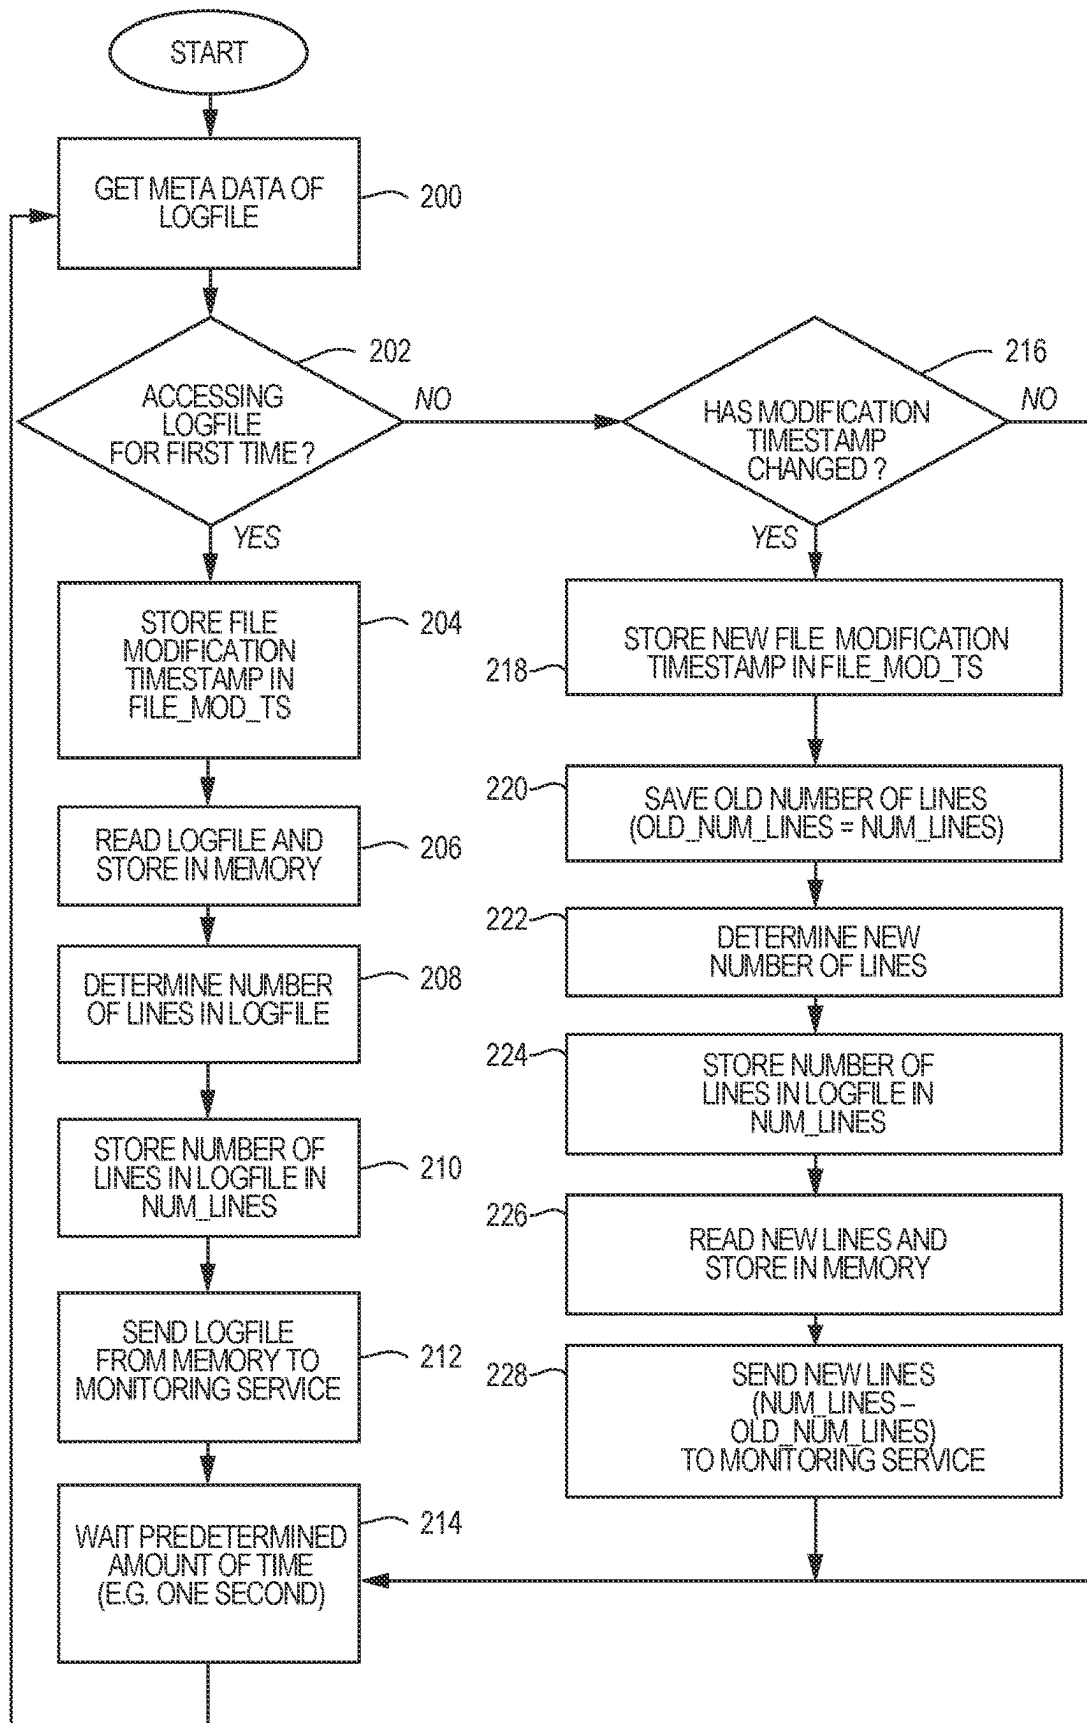


FIG. 4




FIG. 5




FIG. 6

1

LOGFILE COLLECTION AND CONSOLIDATION

RELATED APPLICATION

This application is a continuation of co-pending U.S. patent application Ser. No. 15/202,986, filed on Jul. 6, 2016, entitled "LOGFILE COLLECTION AND CONSOLIDATION," which is hereby incorporated herein by reference in its entirety.

TECHNICAL FIELD

The examples relate to logfiles and, in particular, to logfile collection and consolidation.

BACKGROUND

Computing devices often record events and status information periodically or intermittently that may be subsequently analyzed to glean valuable information. For example, a web server may record each connection made to a website that is hosted on the web server. This connection information may be subsequently analyzed to determine how many connections each website hosted by the web server received over a period of time. Computing devices often record such information in one or more structures referred to as logfiles. In practice, multiple computing devices are often concurrently generating logfiles. For example, an entity may have hundreds of web servers that concurrently host thousands of websites.

SUMMARY

The examples provide mechanisms for logfile collection and consolidation. The examples implement a mechanism by which a plurality of logfiles maintained on a corresponding plurality of remote computing devices can be securely and iteratively consolidated on a remote monitoring service by a single computing device on an ongoing basis.

In one example, a method for consolidating log information from remote computing devices is provided. The method includes establishing, by a computing device comprising a processor device via a network, connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile. The method further includes, for a plurality of iterations, retrieving logfile contents from each logfile on each remote computing device, and sending the logfile contents to a centralized monitoring service.

In one example, for the plurality of iterations, retrieving the logfile contents from each logfile on each remote computing device includes, for each remote computing device, accessing logfile meta data that identifies attributes of the corresponding logfile and determining from the logfile meta data that the logfile has a new logfile status or an old logfile status. If the logfile has the new logfile status, all the logfile contents of the logfile are retrieved. If the logfile has the old logfile status, only the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed are retrieved.

In one example, the logfile meta data comprises inode data. In one example, for each logfile, a first inode identifier that identifies an inode at a time T1 that refers to the respective logfile is stored. A second inode value that identifies the inode at a time T2 that refers to the respective logfile is determined. Responsive to determining that the

2

second inode value is different from the first inode value, the logfile is determined to have the new logfile status.

In another example, a computing device is provided. The computing device includes a memory and a processor device coupled to the memory. Connections with a plurality of remote computing devices are established. Each remote computing device has a corresponding logfile. For a plurality of iterations, logfile contents from each logfile on each remote computing device are retrieved, and the logfile contents are sent to a centralized monitoring service.

Individuals will appreciate the scope of the disclosure and realize additional aspects thereof after reading the following detailed description of the examples in association with the accompanying drawing figures.

BRIEF DESCRIPTION OF THE DRAWINGS

The accompanying drawing figures incorporated in and forming a part of this specification illustrate several aspects of the disclosure and, together with the description, serve to explain the principles of the disclosure.

FIG. 1 is a block diagram of a system in which examples may be practiced;

FIG. 2 is a flowchart of a method for consolidating log information from remote computing devices according to one example;

FIG. 3 is a block diagram of the system illustrated in FIG. 1 that illustrates aspects of a method for consolidating log information from remote computing devices in greater detail;

FIG. 4 is a more detailed flowchart of the method for consolidating log information from the remote computing devices illustrated in FIG. 2;

FIG. 5 is a block diagram of a computing device according to one example; and

FIG. 6 is a block diagram of a simplified version of the system illustrated in FIG. 1 according to one example.

DETAILED DESCRIPTION

The examples set forth below represent the information to enable individuals to practice the examples and illustrate the best mode of practicing the examples. Upon reading the following description in light of the accompanying drawing figures, individuals will understand the concepts of the disclosure and will recognize applications of these concepts not particularly addressed herein. It should be understood that these concepts and applications fall within the scope of the disclosure and the accompanying claims.

Any flowcharts discussed herein are necessarily discussed in some sequence for purposes of illustration, but unless otherwise explicitly indicated, the examples are not limited to any particular sequence of steps. The use herein of ordinals in conjunction with an element is solely for distinguishing what might otherwise be similar or identical labels, such as "first inode identifier" and "second inode identifier," and does not imply a priority, a type, an importance, or other attribute, unless otherwise stated herein.

As used herein and in the claims, the articles "a" and "an" in reference to an element refers to "one or more" of the element unless otherwise explicitly specified.

Computing devices often record events and status information periodically or intermittently that may be subsequently analyzed to glean valuable information. For example, a web server may record each connection made to a website that is hosted on the web server. This connection information may be subsequently analyzed to determine

how many connections each website hosted by the web server received over a period of time. Computing devices often record such information in one or more structures referred to as logfiles. In practice, multiple computing devices are often concurrently generating logfiles. For example, an entity may have hundreds of web servers that concurrently host thousands of websites.

A centralized monitoring service that runs on a monitoring computing device that is different from the computing devices that generate the logfiles is often used to analyze logfiles. The centralized monitoring service typically provides specialized functionality associated with the analysis of logfiles, such as an ability to index key words, build suitable searching structures, and the like. The centralized monitoring service may also offer a user interface that allows an operator or other user to query or otherwise access the logfile contents to glean the desired information.

In order to analyze logfile contents, the logfile contents are first copied from the remote computing devices to the centralized monitoring service. Because the logfile contents in the logfiles are typically generated continually, and because hundreds or even thousands of remote computing devices may be generating logfiles, the logistics of continually moving the logfiles from the remote computing devices to the centralized monitoring service can be time-consuming, or impracticable.

The examples herein establish connections with a plurality of different remote computing devices, each of which has a corresponding logfile. For a plurality of iterations, logfile contents from each logfile on each remote computing device are retrieved, and the logfile contents are sent to a centralized monitoring service. Among other advantages, the examples eliminate a need to manually and repeatedly copy logfiles from multiple different devices to a centralized location, furthermore, the examples only copy logfile contents that have been added since the previous iteration.

FIG. 1 is a block diagram of a system 10 in which examples may be practiced. The system 10 includes a computing device 12 and a plurality of remote computing devices 14-1-14-3 (generally, remote computing devices 14). The phrase "remote" in this context is used solely to differentiate the remote computing devices 14 from the computing device 12, and does not imply a geographic distance or other characteristic or attribute. The remote computing devices 14 could be located in a same building or structure as the computing device 12, or could be located thousands of miles away from the computing device 12. Each remote computing device 14 may be located geographically distant from each other remote computing device 14, or the remote computing devices 14 may be located in relatively close proximity to one another. The computing device 12 communicates with the remote computing devices 14 via one or more networks 16.

The remote computing device 14-1 contains a logfile 18-1. The logfile 18-1 may contain any data that may be periodically and/or intermittently recorded by the remote computing device 14-1. The data may comprise operational aspects of the remote computing device 14-1, such as memory utilization, disk storage utilization, processor utilization, or the like, at a particular instance in time. The data may also comprise events that occur on the remote computing device 14-1, such as detected faults, security breaches, and the like. The nature of the logfile 18-1 may depend on the nature of the application that generates the logfile 18-1. A web server, for example, may record information relating to each request or connection made to a website hosted by the web server. The particular information that is recorded

can comprise any desired information, such as which website received the request, the time of the request, the IP address from which the request originated, and the like. By way of non-limiting example, the information recorded in the logfile 18-1 may comprise requests processed by a server; diagnostic information identifying errors encountered in processing requests; data identifying when a processing task was executed, a name of the processing task, and if the processing task failed or was successful; messages from an operating system kernel that identify problems during the execution of the kernel; and warning or error messages generated by an executing task. The information contained in the logfile 18-1 and other logfiles discussed herein may be referred to as the logfile contents.

In this example the logfile contents comprise a plurality of lines 20-1-20-8 (generally, lines 20). Each line 20 represents a recorded, or logged, record of information. As discussed above, the information may comprise any information that the remote computing device 14-1 is configured to record. The lines 20 are recorded iteratively over time, and thus the logfile 18-1 grows over time. In some environments, after the logfile 18-1 reaches a predetermined size, the logfile 18-1 is automatically closed, renamed, and a new, empty logfile 18-1 is created to prevent a logfile 18-1 from growing beyond a certain size.

The remote computing device 14-1 also includes logfile meta data 22-1 that identifies attributes of the corresponding logfile 18-1, including the locations of the blocks of data that compose the logfile 18-1. The precise content of the logfile meta data 22-1 may differ depending on environmental aspects of the remote computing device 14-1. For example, if the operating system environment of the remote computing device 14-1 is a Unix or Linux operating system, the logfile meta data 22-1 may comprise an inode. The inode data may include, by way of non-limiting example, an inode number that uniquely identifies the respective inode from other inodes, an Access Control List (ACL), extended attributes, pointers to direct/indirect disk blocks, number of blocks, file access time, file creation time, last modification timestamp, file deletion time, file generation number, file size, file type, group, number of links, owner, permissions, and status flags. In other operating environments, such as the Microsoft Windows operating system and the Apple OSX operating system, the logfile meta data 22-1 may have a different format, but generally comprises substantially similar information as discussed herein with regard to the Unix or Linux operating systems.

The remote computing device 14-1 may be addressable by the computing device 12 for purposes of communication via address information 24, which may include, by way of non-limiting example, a hostname, an IP address, or any other identifier via which the computing device 12 may initiate communications with the remote computing device 14-1 via the network 16.

The remote computing device 14-2 similarly contains a logfile 18-2, and logfile meta data 22-2 that identifies attributes of the corresponding logfile 18-2, including the locations of the blocks of data that compose the logfile 18-2. Substantially the same events or environmental information may be logged to the logfile 18-2 as logged to the logfile 18-1. For example, in the example of a web server, the logfile 18-2 may contain a record of information that identifies each request or connection made to a website hosted by the remote computing device 14-2.

The remote computing device 14-3 similarly contains a logfile 18-3, and logfile meta data 22-3 that identifies attributes of the corresponding logfile 18-3, including the

locations of the blocks of data that compose the logfile **18-3**. Substantially the same events or environmental information may be logged to the logfile **18-3** as logged to the logfiles **18-1**, **18-2**. For example, in the example of a web server, the logfile **18-3** may contain a record of information that identifies each request or connection made to a website hosted by the remote computing device **14-3**.

While for purposes of illustration only three remote computing devices **14** are illustrated, the examples are not limited to any particular number of remote computing devices **14**, and have applicability to hundreds or thousands of remote computing devices **14**. Similarly, while for purposes of illustration only, each remote computing device **14** contains only a single logfile **18**, in practice a remote computing device **14** may concurrently maintain any number of logfiles **18**.

The system **10** includes a centralized monitoring service **26** that is implemented on a monitoring computing device **28**. As will be discussed in greater detail herein, the logfile contents of the logfiles **18** are periodically provided to the centralized monitoring service **26**. The centralized monitoring service **26** includes an indexer module **30** that receives the logfile contents of the logfiles **18**, and generates indexed data **32** based on the logfile contents. A full text search engine **34** can search the indexed data **32** in response to queries received from a user who interfaces with the full text search engine **34** via a web user interface **36**.

The computing device **12** includes an agent module **38** that, as discussed in greater detail herein, iteratively and securely collects the logfile contents of the logfiles **18** from the remote computing devices **14**, and sends the logfile contents to the centralized monitoring service **26**. Because the agent module **38** executes on the computing device **12**, functionality implemented by the agent module **38** may be attributed to the computing device **12** throughout the specification and claims herein. The computing device **12** includes a memory **40** that includes configuration information **41** suitable for implementing aspects of the examples. The configuration information **41** may be configured by an operator and maintained in a persistent storage device, and subsequently read or otherwise loaded into the memory **40** by the agent module **38** or other process.

In some examples, the configuration information **41** includes digital key information **42** for facilitating secure connections between the computing device **12** and the remote computing devices **14**. The digital key information **42** comprises a plurality of digital keys **44-1-44-3**, each of which comprises a public key associated with a corresponding remote computing device **14**. In particular, the digital key **44-1** is associated with the remote computing device **14-1**, the digital key **44-2** is associated with the remote computing device **14-2**, and the digital key **44-3** is associated with the remote computing device **14-3**. The digital key information **42** may be used to establish secure connections, such as, by way of non-limiting example, secure shell (SSH) connections, but the examples are not limited to any particular type of secure connections. Such secure connections result in communications between the computing device **12** and the remote computing devices **14** being encrypted. Encrypted communications prevent access to the log contents of the logfiles **18-1-18-3** by any unintended or inadvertent recipients as the log contents of the logfiles **18-1-18-3** are communicated from the remote computing devices **14** to the computing device **12** as discussed in greater detail below.

The configuration information **41** also includes remote log information **46**. The remote log information **46** includes

logfile location records **48-1-48-3**, each of which identifies a logfile to be consolidated, and the location of the respective logfile. The logfile location record **48-1** identifies the logfile **18-1** and provides the hostname of the remote computing device **14-1** so that secure communications between the computing device **12** and the remote computing device **14-1** can be established. The logfile location record **48-2** identifies the logfile **18-2** and provides the hostname of the remote computing device **14-2**. The logfile location record **48-3** identifies the logfile **18-3** and provides the hostname of the remote computing device **14-3**.

FIG. 2 is a flowchart of a method for consolidating log information from the remote computing devices **14** according to one example. FIG. 2 will be discussed in conjunction with FIG. 1. The computing device **12** establishes secure connections with the remote computing devices **14-1** (block **100**). In one example, the secure connections comprise SSH connections. The computing device **12** then retrieves the logfile contents from each logfile **18** on each remote computing device **14** (block **102**). Mechanisms for retrieving the logfile contents are discussed in greater detail below with regard to FIGS. 3 and 4. The computing device **12** then sends the logfile contents to the centralized monitoring service **26** (block **104**). The process described in blocks **102** and **104** may be iteratively performed over a desired period of time, or may continue indefinitely. The iterations may be performed in response to a signal received from another module, in response to a determination that logfile contents have been added to a logfile **18**, or at a desired periodic interval, such as multiple times a second, each second, each desired number of seconds, each minute, or each desired number of minutes.

FIG. 3 is a block diagram of the system **10** illustrated in FIG. 1 that illustrates aspects of a method for consolidating log information from the remote computing devices **14** in greater detail. FIG. 4 is a more detailed flowchart of the method for consolidating log information from the remote computing devices **14** as illustrated in FIG. 2. FIGS. 3 and 4 will be discussed in conjunction with one another. Referring first to FIG. 4, for purposes of clarity and illustration, the method for consolidating log information will be described in detail with respect to a single logfile **18-1** of the remote computing device **14-1**. However, in operation, the same method would be utilized for each logfile **18**. In some examples, the agent module **38** may initiate a separate process, or thread, for each individual logfile **18**, such that the method described herein is performed in parallel for each logfile **18**.

Assume that the computing device **12** has established an SSH connection with the remote computing device **14-1**, as discussed above with regard to block **100** in FIG. 2. At a time period beginning at T1 the computing device **12** obtains logfile meta data **22-1** associated with the logfile **18-1** (block **200**) to determine, among other things, a status of the logfile **18-1**. If the logfile **18-1** has not been previously accessed, the logfile **18-1** has a new logfile status. If the logfile **18-1** has been previously accessed, the logfile **18-1** has an old logfile status. In one example, the logfile meta data **22-1** comprises an inode that contains inode data. To obtain the inode data, the computing device **12** may execute a command via the SSH connection on the remote computing device **14-1**. In one example, the command may comprise the "stat" command. In response to the issuance of the "stat" command, the remote computing device **14-1** provides the inode data to the computing device **12**. The inode data includes information such as the unique inode identifier that uniquely identifies the inode, and the most recent modifi-

cation timestamp associated with the logfile **18-1**. Based on the inode data, the computing device **12** determines whether the logfile **18-1** is being accessed for the first time (block **202**). In particular, in one example, the computing device **12** maintains a memory location **50** in which the inode identifier of the inode is stored each time the inode data is obtained. Because the memory location **50** upon initialization contains no value, the computing device **12** determines that this is the first time the logfile **18-1** has been accessed. The computing device **12** stores the inode identifier from the inode data in the memory location **50**.

The computing device **12** obtains the most recent modification timestamp of the logfile **18-1** from the inode data and stores the most recent modification timestamp in a memory location **52** (block **204**). Assume for purposes of illustration that the modification timestamp at the time **T1** has a modification timestamp value **54-1**. The computing device **12** reads the entire logfile **18-1** from the remote computing device **14-1** and stores the logfile in the memory **40** (block **206**). In one example, the computing device **12** may read the logfile **18-1** by executing a “cat” command via the SSH connection on the remote computing device **14-1**.

The computing device determines the current total number of lines **20** of the logfile **18-1** (block **208**). For purposes of illustration, assume that at time **T1**, the logfile **18-1** contains the four lines **20-1-20-4**. The current total number of lines can be determined in any desired manner. In one example, the Python™ language may be used to implement some of the functionality disclosed herein, and the Python™ “len” function may be used to determine the current total number of lines **20-1-20-4**. The computing device **12** stores the current total number of lines **20-1-20-4** in a memory location **56** (block **210**). Thus, the memory location **56** has a value **58-1**. The computing device **12** sends the lines **20-1-20-4** to the centralized monitoring service **26** (block **212**). The computing device **12** may send the lines **20-1-20-4** to the centralized monitoring service **26** in any desired manner. In one example, the computing device **12** may use the Secure Sockets Layer (SSL) to establish an encrypted link between the computing device **12** and the centralized monitoring service **26**. The first iteration is now complete.

While for purposes of illustration the centralized monitoring service **26** is shown as being implemented on a monitoring computing device **28** that is separate from the computing device **12**, in other examples, the centralized monitoring service **26** may be a component of the computing device **12**, and the computing device **12** sends the lines **20-1-20-4** to the centralized monitoring service **26** via, for example, an interprocess communication mechanism, such as an application programming interface, a port, or the like, or by storing the lines **20-1-20-4** in a storage device to which the centralized monitoring service **26** has access.

The computing device **12** waits a predetermined amount of time (block **214**). As discussed above, the predetermined amount of time may be any desired time interval. After the predetermined amount of time lapses, at a time period beginning at time **T2** the computing device **12** begins the process again and obtains the logfile meta data **22-1** associated with the logfile **18-1** (block **200**). The computing device **12** obtains the inode identifier from the inode data and compares the inode identifier to the previously stored inode identifier stored in the memory location **50** to determine if the logfile **18-1** is being accessed for the first time (block **202**). If the current inode identifier differs from the previously stored inode identifier, the logfile **18-1** is a new logfile, and the logfile **18-1** has the new logfile status. This may occur, for example, if the logfile **18-1** was rotated and

a new logfile **18-1** created since the computing device **12** last accessed the inode data. If the logfile **18-1** is a new logfile, the process described above with regard to blocks **204-214** is repeated.

Assume for purposes of illustration that the current inode identifier matches the inode identifier stored in the memory location **50**. If the current inode identifier matches the inode identifier stored in the memory location **50**, the logfile **18-1** is not being accessed for the first time, and the logfile **18-1** has the old logfile status. The computing device **12** then extracts the most recent modification timestamp from the inode data and compares the most recent modification timestamp with the previous modification timestamp stored in the memory location **52** (block **216**). If the most recent modification timestamp and the previous modification timestamp have not changed, then no additional lines **20** have been added to the logfile **18-1** since the previous iteration, and the process returns to block **214** to wait for the next iteration.

For purposes of illustration, assume that since the last iteration and prior to the time period beginning at time **T2** the remote computing device **14-1** recorded, or logged, the additional lines **20-5-20-6** to the logfile **18-1**. This modification of the logfile **18-1** alters the modification timestamp in the inode data and thus, at block **216**, the computing device **12** determines that the modification timestamp in the inode data differs from the modification timestamp stored in the memory location **52**. The computing device **12** stores the new modification timestamp in the memory location **52**, and the memory location **52** now has the modification timestamp value **54-2** (block **218**). The computing device **12** copies the number of previous lines **20** stored in the memory location **56** to a memory location **59** (block **220**). The memory location **59** has the value **60-1**.

The computing device **12** determines the new current total number of lines **20-1-20-6** in the logfile **18-1** (block **222**), and stores the new current total number of lines **20-1-20-6** in the memory location **56** (block **224**). The memory location **56** has a value **58-2**. The computing device **12** determines the number of new lines **20-5-20-6** based on the current total number of lines **20-1-20-6** and the previous total number of lines **20-1-20-4**. In this example, the number of new lines **20-5-20-6** is two. The computing device **12** reads the number of new lines **20-5-20-6** from the remote computing device **14-1** and stores the new lines **20-5-20-6** in the memory **40** (block **226**). The computing device **12** sends the lines **20-5-20-6** to the centralized monitoring service **26** (block **228**). This completes the second iteration.

The computing device **12** waits a predetermined amount of time (block **214**). After the predetermined amount of time lapses, at a time period beginning at time **T3**, the computing device **12** begins the process again and obtains the logfile meta data **22-1** associated with the logfile **18-1** (block **200**). The computing device **12** obtains the inode identifier from the inode data and compares the inode identifier to the previously stored inode identifier stored in the memory location **50** to determine if the logfile **18-1** is being accessed for the first time (block **202**). Assume for purposes of illustration that the current inode identifier matches the inode identifier stored in the memory location **50**. If the current inode identifier matches the inode identifier stored in the memory location **50**, the logfile **18-1** is not being accessed for the first time. The computing device **12** then extracts the most recent modification timestamp from the inode data and compares the most recent modification timestamp with the previous modification timestamp stored in the memory location **52** (block **216**). For purposes of

illustration, assume that since the last iteration and prior to the time T3 the remote computing device 14-1 recorded, or logged, the additional lines 20-7-20-8 to the logfile 18-1. This modification of the logfile 18-1 alters the modification timestamp in the inode data and thus, at block 216, the computing device 12 determines that the modification timestamp in the inode data differs from the modification timestamp stored in the memory location 52. The computing device 12 stores the new modification timestamp in the memory location 52, and the memory location 52 now has the modification timestamp value 54-3 (block 218). The computing device 12 copies the number of previous lines 20 stored in the memory location 56 to a memory location 59 (block 220). The memory location 59 has the value 60-2.

The computing device 12 determines the new current total number of lines 20-1-20-8 in the logfile 18-1 (block 222), and stores the new current total number of lines in the memory location 56 (block 224). The memory location 56 has a value 58-3. The computing device 12 determines the number of new lines 20-7-20-8 based on the current total number of lines 20-1-20-8 and the previous total number of lines 20-1-20-6. In this example, the number of new lines 20-7-20-8 is two. The computing device 12 reads the number of new lines 20-7-20-8 from the remote computing device 14-1 and stores the new lines 20-7-20-8 in the memory 40 (block 226). The computing device 12 sends the lines 20-7-20-8 to the centralized monitoring service 26 (block 228). This completes the third iteration.

This process may continue indefinitely, and may be performed for the remote computing device 14-2-14-2 consecutively, or in parallel with the remote computing device 14-1. In this manner, the computing device 12 securely consolidates the logfile contents of the logfiles 18-1-18-3 on a continuous basis from the remote computing devices 14-1-14-3 to the centralized monitoring service 26.

While for purposes of illustration the process has been described in terms of a number of lines of a logfile 18, the examples are not limited to using a number of lines of the logfile 18, and can utilize any mechanism by which additional content that has been added to the logfile 18 since the last iteration can be determined. For example, in some examples the computing device 12 may utilize a size of the logfile 18, and copy all logfile content that has been added to a logfile 18 since the previous iteration based on an increase in size of the logfile 18. For example, at each iteration the computing device 12 may determine and store a total size of the logfile 18 in the memory 40. At the successive iteration, the computing device 12 may retrieve the previous size of the logfile 18 from the memory 40, determine a current size of the logfile 18, determine a new amount of data based on the previous size and the current size, and retrieve a most recent amount of data added to the logfile 18 equal to the new amount of data.

For example, during one iteration a logfile 18 may be 434 KB in size. At a next iteration, the logfile 18 may have grown to 437 KB in size, and the computing device 12 may read the last 3 kB of data from the logfile 18 and send the data to the centralized monitoring service 26.

FIG. 5 is a block diagram of the computing device 12 suitable for implementing examples according to one example. The computing device 12 may comprise any computing or electronic device capable of including firmware, hardware, and/or executing software instructions to implement the functionality described herein, such as a computer server device, a desktop computing device, or the like. The computing device 12 includes a processor device 62, the system memory 40, and a system bus 64. The system

bus 64 provides an interface for system components including, but not limited to, the system memory 40 and the processor device 62. The processor device 62 can be any commercially available or proprietary processor.

The system bus 64 may be any of several types of bus structures that may further interconnect to a memory bus (with or without a memory controller), a peripheral bus, and/or a local bus using any of a variety of commercially available bus architectures. The system memory 40 may include non-volatile memory 66 (e.g., read-only memory (ROM), erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), etc.), and/or volatile memory 68 (e.g., random-access memory (RAM)). A basic input/output system (BIOS) 70 may be stored in the non-volatile memory 66 and can include the basic routines that help to transfer information between elements within the computing device 12. The volatile memory 68 may also include a high-speed RAM, such as static RAM, for caching data.

The computing device 12 may further include or be coupled to a non-transitory computer-readable storage medium or storage device 72, which may comprise, for example, an internal or external hard disk drive (HDD) (e.g., enhanced integrated drive electronics (EIDE) or serial advanced technology attachment (SATA)), HDD (e.g., EIDE or SATA) for storage, flash memory, or the like. The storage device 72 and other drives associated with computer-readable media and computer-usable media may provide non-volatile storage of data, data structures, computer-executable instructions, and the like, including, for example, the configuration information 41. Although the description of computer-readable media above refers to an HDD, it should be appreciated that other types of media that are readable by a computer, such as Zip disks, magnetic cassettes, flash memory cards, cartridges, and the like, may also be used in the operating environment, and, further, that any such media may contain computer-executable instructions for performing novel methods of the disclosed examples.

A number of modules can be stored in the storage device 72 and in the volatile memory 68, including an operating system 74 and one or more program modules 76, which may implement the functionality described herein in whole or in part, including, for example, the functionality described herein with regard to the agent module 38. It is to be appreciated that the examples can be implemented with various commercially available operating systems 74.

All or a portion of the examples may be implemented as a computer program product stored on a transitory or non-transitory computer-usable or computer-readable storage medium, such as the storage device 72, which includes complex programming instructions, such as complex computer-readable program code, configured to cause the processor device 62 to carry out the steps described herein. Thus, the computer-readable program code can comprise software instructions for implementing the functionality of the examples described herein when executed on the processor device 62. The processor device 62, in conjunction with the program modules 76 in the volatile memory 68, may serve as a controller, or control system, for the computing device 12 that is configured to, or adapted to, implement the functionality described herein.

An operator or other user may also be able to enter one or more configuration commands through a keyboard (not illustrated), a pointing device such as a mouse (not illustrated), or a touch-sensitive surface (not illustrated). Such input devices may be connected to the processor device 62 through an input device interface 78 that is coupled to the

11

system bus **64** but can be connected by other interfaces such as a parallel port, an Institute of Electrical and Electronic Engineers (IEEE) 1394 serial port, a Universal Serial Bus (USB) port, an IR interface, and the like.

The computing device **12** also includes a communication interface **80** suitable for communicating with the network **16** as appropriate or desired. The computing device **12** may also include a video port **82** configured to interface with a display **84**, to provide the operator or user information during the examples disclosed herein.

FIG. **6** is a block diagram of a simplified version of the system **10** illustrated in FIG. **1**, according to one example. The system **10** includes the computing device **12** and the plurality of remote computing devices **14**. The computing device **12** includes the memory **40** and the processor device **62** coupled to the memory **40**. The processor device **62** establishes connections with the plurality of remote computing devices **14**. Each remote computing device **14** has a corresponding logfile **18**. For a plurality of iterations, the processor device **62** retrieves logfile contents from each logfile **18** on each remote computing device **14**, and sends the logfile contents to the centralized monitoring service **26**.

Individuals will recognize improvements and modifications to the examples of the disclosure. All such improvements and modifications are considered within the scope of the concepts disclosed herein and the claims that follow.

What is claimed is:

1. A method for consolidating logfile contents from remote computing devices, comprising:

establishing, by a computing device comprising a processor device via a network, connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile comprising information that the remote computing device is configured to store in the corresponding logfile; and for a plurality of iterations, by the computing device via the network:

for each remote computing device, accessing logfile meta data that identifies attributes of the corresponding logfile, the logfile meta data being separate from data stored in the corresponding logfile, wherein the logfile meta data comprises inode data indicative of whether the corresponding logfile was previously accessed; and

for each remote computing device, based on the inode data indicative of whether the corresponding logfile was accessed, retrieving at least a portion of logfile contents from the corresponding logfile, wherein retrieving the at least the portion of the logfile contents comprises:

determining whether a first inode value of the inode data that identifies an inode at a time T1 and that refers to the corresponding logfile is different from a second inode value of the inode data that identifies the inode at a time T2;

based on the first inode value being different from the second inode value, determining that the logfile was previously accessed and that the logfile has an old logfile status;

retrieving the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed; and

sending the at least the portion of the logfile contents to a centralized monitoring service.

2. The method of claim **1** wherein retrieving the at least the portion of the logfile contents further comprises deter-

12

mining a total number of lines of the logfile and storing the total number of lines in a memory.

3. The method of claim **2** wherein retrieving the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed comprises:

retrieving a previous total number of lines in the logfile from the memory;

determining a current total number of lines in the logfile; determining a number of new lines in the logfile based on the previous total number of lines in the logfile and the current total number of lines in the logfile; and

retrieving a number of most recent lines written to the logfile equal to the number of new lines.

4. The method of claim **1** further comprising:

storing, for each logfile, the first inode value that identifies the inode at the time T1 that refers to the corresponding logfile.

5. The method of claim **1** wherein retrieving the at least the portion of the logfile contents further comprises determining a total size of the logfile and storing the total size of the logfile in a memory.

6. The method of claim **5** wherein retrieving the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed comprises:

retrieving a previous size of the logfile from the memory; determining a current size of the logfile;

determining a new amount of data based on the previous size and the current size; and

retrieving a most recent amount of data added to the logfile equal to the new amount of data.

7. The method of claim **1** further comprising waiting a predetermined amount of time, and after the predetermined amount of time, performing another iteration of the plurality of iterations.

8. The method of claim **1** wherein establishing the connections with the plurality of remote computing devices comprises establishing secure connections with the plurality of remote computing devices.

9. The method of claim **8** wherein establishing secure connections with the plurality of remote computing devices comprises obtaining, for each remote computing device of the plurality of remote computing devices, a public key associated with the remote computing device, and establishing the secure connections with the plurality of remote computing devices using the public key associated with each respective remote computing device.

10. A computing device comprising:

a memory; and

a processor device coupled to the memory to:

establish connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile comprising information that the remote computing device is configured to store in the corresponding logfile; and

for a plurality of iterations:

for each remote computing device, access logfile meta data that identifies attributes of the corresponding logfile, the logfile meta data being separate from data stored in the corresponding logfile, wherein the logfile meta data comprises inode data indicative of whether the corresponding logfile was previously accessed; and

for each remote computing device, based on the inode data indicative of whether the corresponding logfile was accessed, retrieve at least a portion

13

of logfile contents from the corresponding logfile, and wherein, to retrieve the at least the portion of the logfile contents, the processor device is further to:

determine whether a first inode value of the inode data that identifies an inode at a time T1 and that refers to the corresponding logfile is different from a second inode value of the inode data that identifies the inode at a time T2;

based on the first inode value being different from the second inode value, determine that the logfile was previously accessed and that the logfile has an old logfile status;

retrieve the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed; and

send the at least the portion of the logfile contents to a centralized monitoring service.

11. The computing device of claim 10 wherein, to retrieve the at least the portion of the logfile contents, the processor device is further to determine a total number of lines of the logfile, and store the total number of lines in a memory.

12. The computing device of claim 11 wherein to retrieve the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed the processor device is further to:

retrieve a previous total number of lines in the logfile from the memory;

determine a current total number of lines in the logfile;

determine a number of new lines in the logfile based on the previous total number of lines in the logfile and the current total number of lines in the logfile; and

retrieve a number of most recent lines written to the logfile equal to the number of new lines.

13. A computer program product for consolidating log information from remote computing devices, the computer program product stored on a non-transitory computer-read-

14

able storage medium and including instructions to cause a processor device to carry out steps of:

establishing connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile comprising information that the remote computing device is configured to store in the corresponding logfile;

for a plurality of iterations, by a computing device via a network:

for each remote computing device, accessing logfile meta data that identifies attributes of the corresponding logfile, the logfile meta data being separate from data stored in the corresponding logfile, wherein the logfile meta data comprises inode data indicative of whether the corresponding logfile was previously accessed; and

for each remote computing device, based on the inode data indicative of whether the corresponding logfile was accessed, retrieving at least a portion of logfile contents from the corresponding logfile, and wherein retrieving the at least the portion of the logfile contents comprises:

determining whether a first inode value of the inode data that identifies an inode at a time T1 and that refers to the corresponding logfile is different from a second inode value of the inode data that identifies the inode at a time T2;

based on the first inode value being different from the second inode value, determining that the logfile was previously accessed and that the logfile has an old logfile status;

retrieving the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed; and

sending the at least the portion of the logfile contents to a centralized monitoring service.

* * * * *