

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393567
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Shuster; Boaz

Logfile collection and consolidation

Abstract

Mechanisms for consolidating log information from remote computing devices are provided. Connections with a plurality of remote computing devices are established. Each remote computing device has a corresponding logfile. For a plurality of iterations, logfile contents from each logfile on each remote computing device are retrieved, and the logfile contents are sent to a centralized monitoring service.

Inventors:	Shuster; Boaz (Kfar Saba, IL)
Applicant:	Red Hat, Inc. (Raleigh, NC)
Family ID:	1000008765698
Assignee:	Red Hat, Inc. (Raleigh, NC)
Appl. No.:	17/532023
Filed:	November 22, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220075769 A1	Mar. 10, 2022

Related U.S. Application Data

continuation parent-doc US 15202986 20160706 US 11188513 child-doc US 17532023

Publication Classification

Int. Cl.: G06F16/23 (20190101); G06F16/18 (20190101); G06F16/182 (20190101);
H04L43/0817 (20220101)

U.S. Cl.:

CPC **G06F16/2322** (20190101); **G06F16/1805** (20190101); **G06F16/2358** (20190101);
G06F16/182 (20190101); H04L43/0817 (20130101)

Field of Classification Search

CPC: G06F (16/2322); G06F (16/2358); G06F (16/1805); G06F (16/182)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5978940	12/1998	Newman et al.	N/A	N/A
5987633	12/1998	Newman et al.	N/A	N/A
6374296	12/2001	Lim et al.	N/A	N/A
6789156	12/2003	Waldspurger	N/A	N/A
7209898	12/2006	Pfeiffer et al.	N/A	N/A
7653836	12/2009	Chatterjee et al.	N/A	N/A
8464335	12/2012	Sinha et al.	N/A	N/A
9736243	12/2016	Chablani et al.	N/A	N/A
9959489	12/2017	Yamamoto et al.	N/A	N/A
10142195	12/2017	Sledz et al.	N/A	N/A
10181991	12/2018	Mathews et al.	N/A	N/A
10387370	12/2018	Shuster	N/A	N/A
2003/0009486	12/2002	Shimada et al.	N/A	N/A
2003/0105811	12/2002	Laborde et al.	N/A	N/A
2003/0135507	12/2002	Hind et al.	N/A	N/A
2004/0078722	12/2003	Pfeiffer et al.	N/A	N/A
2005/0235058	12/2004	Rackus et al.	N/A	N/A
2005/0273667	12/2004	Shrivastava et al.	N/A	N/A
2006/0106851	12/2005	Warshawsky	N/A	G06F 16/2358
2007/0050777	12/2006	Hutchinson et al.	N/A	N/A
2007/0100991	12/2006	Daniels et al.	N/A	N/A
2007/0283194	12/2006	Villella	714/E11.204	G06F 11/3476
2009/0013149	12/2008	Uhlig et al.	N/A	N/A
2009/0138703	12/2008	Schneider	N/A	N/A
2009/0313503	12/2008	Atluri et al.	N/A	N/A
2011/0153616	12/2010	Torres	N/A	N/A
2011/0208663	12/2010	Kennis	705/317	G06Q 30/018
2011/0276695	12/2010	Maldaner	N/A	N/A
2012/0017028	12/2011	Tsirkin	N/A	N/A
2012/0216007	12/2011	Tsirkin	711/E12.002	G06F 9/526
2013/0227352	12/2012	Kumarasamy et al.	N/A	N/A
2013/0332424	12/2012	Nos	707/E17.014	G06F 16/1734
2014/0040180	12/2013	Ruhl et al.	N/A	N/A
2014/0058801	12/2013	Deodhar et al.	N/A	N/A
2014/0181036	12/2013	Dhamankar	707/648	G06F 16/1734
2014/0181461	12/2013	Kegel et al.	N/A	N/A

2014/0279920	12/2013	Madhavarapu	707/649	G06F 16/2358
2014/0282031	12/2013	Hinterbichler et al.	N/A	N/A
2015/0220605	12/2014	Syed et al.	N/A	N/A
2015/0256423	12/2014	Stearns	N/A	N/A
2015/0269183	12/2014	Avati	707/634	G06F 16/1865
2015/0269213	12/2014	Avati	707/693	G06F 16/9017
2015/0286663	12/2014	Mangan, III et al.	N/A	N/A
2015/0324297	12/2014	Tati et al.	N/A	N/A
2016/0041976	12/2015	Pradeep	707/695	G06F 16/2358
2016/0275101	12/2015	Hinterbichler	N/A	G06F 16/1734
2016/0275128	12/2015	Dayka	N/A	G06F 16/2358
2016/0292263	12/2015	Ferrar	N/A	N/A
2016/0299934	12/2015	Karandikar	N/A	G06F 16/248
2016/0378577	12/2015	Russell et al.	N/A	N/A
2017/0075922	12/2016	Torman	N/A	G06F 16/1805
2017/0104820	12/2016	Golander et al.	N/A	N/A
2017/0331880	12/2016	Crofton	N/A	G06F 16/182
2018/0013612	12/2017	Shuster	N/A	N/A

OTHER PUBLICATIONS

Non-Final Office Action for U.S. Appl. No. 15/202,986, mailed Jun. 1, 2019, 23 pages. cited by applicant

Final Office Action for U.S. Appl. No. 15/202,986, mailed Jan. 18, 2019, 18 pages. cited by applicant

Advisory Action for U.S. Appl. No. 15/202,986, mailed May 23, 2019, 5 pages. cited by applicant

Non-Final Office Action for U.S. Appl. No. 15/202,986, mailed Sep. 17, 2019, 20 pages. cited by applicant

Final Office Action for U.S. Appl. No. 15/202,986, mailed Apr. 22, 2020, 25 pages. cited by applicant

Advisory Action for U.S. Appl. No. 15/202,986, mailed Jul. 17, 2020, 3 pages. cited by applicant

Non-Final Office Action for U.S. Appl. No. 15/202,986, mailed Jan. 14, 2021, 26 pages. cited by applicant

Notice of Allowance for U.S. Appl. No. 15/202,986, mailed Jul. 30, 2021, 13 pages. cited by applicant

Primary Examiner: Singh; Amresh

Attorney, Agent or Firm: Dority & Manning, P.A.

Background/Summary

RELATED APPLICATION (1) This application is a continuation of co-pending U.S. patent application Ser. No. 15/202,986, filed on Jul. 6, 2016, entitled “LOGFILE COLLECTION AND CONSOLIDATION,” which is hereby incorporated herein by reference in its entirety.

TECHNICAL FIELD

(1) The examples relate to logfiles and, in particular, to logfile collection and consolidation.

BACKGROUND

(2) Computing devices often record events and status information periodically or intermittently that

may be subsequently analyzed to glean valuable information. For example, a web server may record each connection made to a website that is hosted on the web server. This connection information may be subsequently analyzed to determine how many connections each website hosted by the web server received over a period of time. Computing devices often record such information in one or more structures referred to as logfiles. In practice, multiple computing devices are often concurrently generating logfiles. For example, an entity may have hundreds of web servers that concurrently host thousands of websites.

SUMMARY

(3) The examples provide mechanisms for logfile collection and consolidation. The examples implement a mechanism by which a plurality of logfiles maintained on a corresponding plurality of remote computing devices can be securely and iteratively consolidated on a remote monitoring service by a single computing device on an ongoing basis.

(4) In one example, a method for consolidating log information from remote computing devices is provided. The method includes establishing, by a computing device comprising a processor device via a network, connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile. The method further includes, for a plurality of iterations, retrieving logfile contents from each logfile on each remote computing device, and sending the logfile contents to a centralized monitoring service.

(5) In one example, for the plurality of iterations, retrieving the logfile contents from each logfile on each remote computing device includes, for each remote computing device, accessing logfile meta data that identifies attributes of the corresponding logfile and determining from the logfile meta data that the logfile has a new logfile status or an old logfile status. If the logfile has the new logfile status, all the logfile contents of the logfile are retrieved. If the logfile has the old logfile status, only the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed are retrieved.

(6) In one example, the logfile meta data comprises inode data. In one example, for each logfile, a first inode identifier that identifies an inode at a time T1 that refers to the respective logfile is stored. A second inode value that identifies the inode at a time T2 that refers to the respective logfile is determined. Responsive to determining that the second inode value is different from the first inode value, the logfile is determined to have the new logfile status.

(7) In another example, a computing device is provided. The computing device includes a memory and a processor device coupled to the memory. Connections with a plurality of remote computing devices are established. Each remote computing device has a corresponding logfile. For a plurality of iterations, logfile contents from each logfile on each remote computing device are retrieved, and the logfile contents are sent to a centralized monitoring service.

(8) Individuals will appreciate the scope of the disclosure and realize additional aspects thereof after reading the following detailed description of the examples in association with the accompanying drawing figures.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The accompanying drawing figures incorporated in and forming a part of this specification illustrate several aspects of the disclosure and, together with the description, serve to explain the principles of the disclosure.

(2) FIG. 1 is a block diagram of a system in which examples may be practiced;

(3) FIG. 2 is a flowchart of a method for consolidating log information from remote computing devices according to one example;

(4) FIG. 3 is a block diagram of the system illustrated in FIG. 1 that illustrates aspects of a method

- for consolidating log information from remote computing devices in greater detail;
- (5) FIG. 4 is a more detailed flowchart of the method for consolidating log information from the remote computing devices illustrated in FIG. 2;
- (6) FIG. 5 is a block diagram of a computing device according to one example; and
- (7) FIG. 6 is a block diagram of a simplified version of the system illustrated in FIG. 1 according to one example.

DETAILED DESCRIPTION

- (8) The examples set forth below represent the information to enable individuals to practice the examples and illustrate the best mode of practicing the examples. Upon reading the following description in light of the accompanying drawing figures, individuals will understand the concepts of the disclosure and will recognize applications of these concepts not particularly addressed herein. It should be understood that these concepts and applications fall within the scope of the disclosure and the accompanying claims.
- (9) Any flowcharts discussed herein are necessarily discussed in some sequence for purposes of illustration, but unless otherwise explicitly indicated, the examples are not limited to any particular sequence of steps. The use herein of ordinals in conjunction with an element is solely for distinguishing what might otherwise be similar or identical labels, such as “first inode identifier” and “second inode identifier,” and does not imply a priority, a type, an importance, or other attribute, unless otherwise stated herein.
- (10) As used herein and in the claims, the articles “a” and “an” in reference to an element refers to “one or more” of the element unless otherwise explicitly specified.
- (11) Computing devices often record events and status information periodically or intermittently that may be subsequently analyzed to glean valuable information. For example, a web server may record each connection made to a website that is hosted on the web server. This connection information may be subsequently analyzed to determine how many connections each website hosted by the web server received over a period of time. Computing devices often record such information in one or more structures referred to as logfiles. In practice, multiple computing devices are often concurrently generating logfiles. For example, an entity may have hundreds of web servers that concurrently host thousands of websites.
- (12) A centralized monitoring service that runs on a monitoring computing device that is different from the computing devices that generate the logfiles is often used to analyze logfiles. The centralized monitoring service typically provides specialized functionality associated with the analysis of logfiles, such as an ability to index key words, build suitable searching structures, and the like. The centralized monitoring service may also offer a user interface that allows an operator or other user to query or otherwise access the logfile contents to glean the desired information.
- (13) In order to analyze logfile contents, the logfile contents are first copied from the remote computing devices to the centralized monitoring service. Because the logfile contents in the logfiles are typically generated continually, and because hundreds or even thousands of remote computing devices may be generating logfiles, the logistics of continually moving the logfiles from the remote computing devices to the centralized monitoring service can be time-consuming, or impracticable.
- (14) The examples herein establish connections with a plurality of different remote computing devices, each of which has a corresponding logfile. For a plurality of iterations, logfile contents from each logfile on each remote computing device are retrieved, and the logfile contents are sent to a centralized monitoring service. Among other advantages, the examples eliminate a need to manually and repeatedly copy logfiles from multiple different devices to a centralized location, furthermore, the examples only copy logfile contents that have been added since the previous iteration.
- (15) FIG. 1 is a block diagram of a system 10 in which examples may be practiced. The system 10 includes a computing device 12 and a plurality of remote computing devices 14-1-14-3 (generally, remote computing devices 14). The phrase “remote” in this context is used solely to differentiate

the remote computing devices **14** from the computing device **12**, and does not imply a geographic distance or other characteristic or attribute. The remote computing devices **14** could be located in a same building or structure as the computing device **12**, or could be located thousands of miles away from the computing device **12**. Each remote computing device **14** may be located geographically distant from each other remote computing device **14**, or the remote computing devices **14** may be located in relatively close proximity to one another. The computing device **12** communicates with the remote computing devices **14** via one or more networks **16**.

(16) The remote computing device **14-1** contains a logfile **18-1**. The logfile **18-1** may contain any data that may be periodically and/or intermittently recorded by the remote computing device **14-1**. The data may comprise operational aspects of the remote computing device **14-1**, such as memory utilization, disk storage utilization, processor utilization, or the like, at a particular instance in time. The data may also comprise events that occur on the remote computing device **14-1**, such as detected faults, security breaches, and the like. The nature of the logfile **18-1** may depend on the nature of the application that generates the logfile **18-1**. A web server, for example, may record information relating to each request or connection made to a website hosted by the web server. The particular information that is recorded can comprise any desired information, such as which website received the request, the time of the request, the IP address from which the request originated, and the like. By way of non-limiting example, the information recorded in the logfile **18-1** may comprise requests processed by a server, diagnostic information identifying errors encountered in processing requests; data identifying when a processing task was executed, a name of the processing task, and if the processing task failed or was successful; messages from an operating system kernel that identify problems during the execution of the kernel; and warning or error messages generated by an executing task. The information contained in the logfile **18-1** and other logfiles discussed herein may be referred to as the logfile contents.

(17) In this example the logfile contents comprise a plurality of lines **20-1-20-8** (generally, lines **20**). Each line **20** represents a recorded, or logged, record of information. As discussed above, the information may comprise any information that the remote computing device **14-1** is configured to record. The lines **20** are recorded iteratively over time, and thus the logfile **18-1** grows over time. In some environments, after the logfile **18-1** reaches a predetermined size, the logfile **18-1** is automatically closed, renamed, and a new, empty logfile **18-1** is created to prevent a logfile **18-1** from growing beyond a certain size.

(18) The remote computing device **14-1** also includes logfile meta data **22-1** that identifies attributes of the corresponding logfile **18-1**, including the locations of the blocks of data that compose the logfile **18-1**. The precise content of the logfile meta data **22-1** may differ depending on environmental aspects of the remote computing device **14-1**. For example, if the operating system environment of the remote computing device **14-1** is a Unix or Linux operating system, the logfile meta data **22-1** may comprise an inode. The inode data may include, by way of non-limiting example, an inode number that uniquely identifies the respective inode from other inodes, an Access Control List (ACL), extended attributes, pointers to direct/indirect disk blocks, number of blocks, file access time, file creation time, last modification timestamp, file deletion time, file generation number, file size, file type, group, number of links, owner, permissions, and status flags. In other operating environments, such as the Microsoft Windows operating system and the Apple OSX operating system, the logfile meta data **22-1** may have a different format, but generally comprises substantially similar information as discussed herein with regard to the Unix or Linux operating systems.

(19) The remote computing device **14-1** may be addressable by the computing device **12** for purposes of communication via address information **24**, which may include, by way of non-limiting example, a hostname, an IP address, or any other identifier via which the computing device **12** may initiate communications with the remote computing device **14-1** via the network **16**.

(20) The remote computing device **14-2** similarly contains a logfile **18-2**, and logfile meta data **22-**

2 that identifies attributes of the corresponding logfile **18-2**, including the locations of the blocks of data that compose the logfile **18-2**. Substantially the same events or environmental information may be logged to the logfile **18-2** as logged to the logfile **18-1**. For example, in the example of a web server, the logfile **18-2** may contain a record of information that identifies each request or connection made to a website hosted by the remote computing device **14-2**.

(21) The remote computing device **14-3** similarly contains a logfile **18-3**, and logfile meta data **22-3** that identifies attributes of the corresponding logfile **18-3**, including the locations of the blocks of data that compose the logfile **18-3**. Substantially the same events or environmental information may be logged to the logfile **18-3** as logged to the logfiles **18-1**, **18-2**. For example, in the example of a web server, the logfile **18-3** may contain a record of information that identifies each request or connection made to a website hosted by the remote computing device **14-3**.

(22) While for purposes of illustration only three remote computing devices **14** are illustrated, the examples are not limited to any particular number of remote computing devices **14**, and have applicability to hundreds or thousands of remote computing devices **14**. Similarly, while for purposes of illustration only, each remote computing device **14** contains only a single logfile **18**, in practice a remote computing device **14** may concurrently maintain any number of logfiles **18**.

(23) The system **10** includes a centralized monitoring service **26** that is implemented on a monitoring computing device **28**. As will be discussed in greater detail herein, the logfile contents of the logfiles **18** are periodically provided to the centralized monitoring service **26**. The centralized monitoring service **26** includes an indexer module **30** that receives the logfile contents of the logfiles **18**, and generates indexed data **32** based on the logfile contents. A full text search engine **34** can search the indexed data **32** in response to queries received from a user who interfaces with the full text search engine **34** via a web user interface **36**.

(24) The computing device **12** includes an agent module **38** that, as discussed in greater detail herein, iteratively and securely collects the logfile contents of the logfiles **18** from the remote computing devices **14**, and sends the logfile contents to the centralized monitoring service **26**. Because the agent module **38** executes on the computing device **12**, functionality implemented by the agent module **38** may be attributed to the computing device **12** throughout the specification and claims herein. The computing device **12** includes a memory **40** that includes configuration information **41** suitable for implementing aspects of the examples. The configuration information **41** may be configured by an operator and maintained in a persistent storage device, and subsequently read or otherwise loaded into the memory **40** by the agent module **38** or other process.

(25) In some examples, the configuration information **41** includes digital key information **42** for facilitating secure connections between the computing device **12** and the remote computing devices **14**. The digital key information **42** comprises a plurality of digital keys **44-1-44-3**, each of which comprises a public key associated with a corresponding remote computing device **14**. In particular, the digital key **44-1** is associated with the remote computing device **14-1**, the digital key **44-2** is associated with the remote computing device **14-2**, and the digital key **44-3** is associated with the remote computing device **14-3**. The digital key information **42** may be used to establish secure connections, such as, by way of non-limiting example, secure shell (SSH) connections, but the examples are not limited to any particular type of secure connections. Such secure connections result in communications between the computing device **12** and the remote computing devices **14** being encrypted. Encrypted communications prevent access to the log contents of the logfiles **18-1-18-3** by any unintended or inadvertent recipients as the log contents of the logfiles **18-1-18-3** are communicated from the remote computing devices **14** to the computing device **12** as discussed in greater detail below.

(26) The configuration information **41** also includes remote log information **46**. The remote log information **46** includes logfile location records **48-1-48-3**, each of which identifies a logfile to be consolidated, and the location of the respective logfile. The logfile location record **48-1** identifies

the logfile **18-1** and provides the hostname of the remote computing device **14-1** so that secure communications between the computing device **12** and the remote computing device **14-1** can be established. The logfile location record **48-2** identifies the logfile **18-2** and provides the hostname of the remote computing device **14-2**. The logfile location record **48-3** identifies the logfile **18-3** and provides the hostname of the remote computing device **14-3**.

(27) FIG. 2 is a flowchart of a method for consolidating log information from the remote computing devices **14** according to one example. FIG. 2 will be discussed in conjunction with FIG. 1. The computing device **12** establishes secure connections with the remote computing devices **14-1** (block **100**). In one example, the secure connections comprise SSH connections. The computing device **12** then retrieves the logfile contents from each logfile **18** on each remote computing device **14** (block **102**). Mechanisms for retrieving the logfile contents are discussed in greater detail below with regard to FIGS. 3 and 4. The computing device **12** then sends the logfile contents to the centralized monitoring service **26** (block **104**). The process described in blocks **102** and **104** may be iteratively performed over a desired period of time, or may continue indefinitely. The iterations may be performed in response to a signal received from another module, in response to a determination that logfile contents have been added to a logfile **18**, or at a desired periodic interval, such as multiple times a second, each second, each desired number of seconds, each minute, or each desired number of minutes.

(28) FIG. 3 is a block diagram of the system **10** illustrated in FIG. 1 that illustrates aspects of a method for consolidating log information from the remote computing devices **14** in greater detail. FIG. 4 is a more detailed flowchart of the method for consolidating log information from the remote computing devices **14** as illustrated in FIG. 2. FIGS. 3 and 4 will be discussed in conjunction with one another. Referring first to FIG. 4, for purposes of clarity and illustration, the method for consolidating log information will be described in detail with respect to a single logfile **18-1** of the remote computing device **14-1**. However, in operation, the same method would be utilized for each logfile **18**. In some examples, the agent module **38** may initiate a separate process, or thread, for each individual logfile **18**, such that the method described herein is performed in parallel for each logfile **18**.

(29) Assume that the computing device **12** has established an SSH connection with the remote computing device **14-1**, as discussed above with regard to block **100** in FIG. 2. At a time period beginning at T1 the computing device **12** obtains logfile meta data **22-1** associated with the logfile **18-1** (block **200**) to determine, among other things, a status of the logfile **18-1**. If the logfile **18-1** has not been previously accessed, the logfile **18-1** has a new logfile status. If the logfile **18-1** has been previously accessed, the logfile **18-1** has an old logfile status. In one example, the logfile meta data **22-1** comprises an inode that contains inode data. To obtain the inode data, the computing device **12** may execute a command via the SSH connection on the remote computing device **14-1**. In one example, the command may comprise the “stat” command. In response to the issuance of the “stat” command, the remote computing device **14-1** provides the inode data to the computing device **12**. The inode data includes information such as the unique inode identifier that uniquely identifies the inode, and the most recent modification timestamp associated with the logfile **18-1**. Based on the inode data, the computing device **12** determines whether the logfile **18-1** is being accessed for the first time (block **202**). In particular, in one example, the computing device **12** maintains a memory location **50** in which the inode identifier of the inode is stored each time the inode data is obtained. Because the memory location **50** upon initialization contains no value, the computing device **12** determines that this is the first time the logfile **18-1** has been accessed. The computing device **12** stores the inode identifier from the inode data in the memory location **50**.

(30) The computing device **12** obtains the most recent modification timestamp of the logfile **18-1** from the inode data and stores the most recent modification timestamp in a memory location **52** (block **204**). Assume for purposes of illustration that the modification timestamp at the time T1 has a modification timestamp value **54-1**. The computing device **12** reads the entire logfile **18-1** from

the remote computing device **14-1** and stores the logfile in the memory **40** (block **206**). In one example, the computing device **12** may read the logfile **18-1** by executing a “cat” command via the SSH connection on the remote computing device **14-1**.

(31) The computing device determines the current total number of lines **20** of the logfile **18-1** (block **208**). For purposes of illustration, assume that at time T1, the logfile **18-1** contains the four lines **20-1-20-4**. The current total number of lines can be determined in any desired manner. In one example, the Python™ language may be used to implement some of the functionality disclosed herein, and the Python™ “len” function may be used to determine the current total number of lines **20-1-20-4**. The computing device **12** stores the current total number of lines **20-1-20-4** in a memory location **56** (block **210**). Thus, the memory location **56** has a value **58-1**. The computing device **12** sends the lines **20-1-20-4** to the centralized monitoring service **26** (block **212**). The computing device **12** may send the lines **20-1-20-4** to the centralized monitoring service **26** in any desired manner. In one example, the computing device **12** may use the Secure Sockets Layer (SSL) to establish an encrypted link between the computing device **12** and the centralized monitoring service **26**. The first iteration is now complete.

(32) While for purposes of illustration the centralized monitoring service **26** is shown as being implemented on a monitoring computing device **28** that is separate from the computing device **12**, in other examples, the centralized monitoring service **26** may be a component of the computing device **12**, and the computing device **12** sends the lines **20-1-20-4** to the centralized monitoring service **26** via, for example, an interprocess communication mechanism, such as an application programming interface, a port, or the like, or by storing the lines **20-1-20-4** in a storage device to which the centralized monitoring service **26** has access.

(33) The computing device **12** waits a predetermined amount of time (block **214**). As discussed above, the predetermined amount of time may be any desired time interval. After the predetermined amount of time lapses, at a time period beginning at time T2 the computing device **12** begins the process again and obtains the logfile meta data **22-1** associated with the logfile **18-1** (block **200**). The computing device **12** obtains the inode identifier from the inode data and compares the inode identifier to the previously stored inode identifier stored in the memory location **50** to determine if the logfile **18-1** is being accessed for the first time (block **202**). If the current inode identifier differs from the previously stored inode identifier, the logfile **18-1** is a new logfile, and the logfile **18-1** has the new logfile status. This may occur, for example, if the logfile **18-1** was rotated and a new logfile **18-1** created since the computing device **12** last accessed the inode data. If the logfile **18-1** is a new logfile, the process described above with regard to blocks **204-214** is repeated.

(34) Assume for purposes of illustration that the current inode identifier matches the inode identifier stored in the memory location **50**. If the current inode identifier matches the inode identifier stored in the memory location **50**, the logfile **18-1** is not being accessed for the first time, and the logfile **18-1** has the old logfile status. The computing device **12** then extracts the most recent modification timestamp from the inode data and compares the most recent modification timestamp with the previous modification timestamp stored in the memory location **52** (block **216**). If the most recent modification timestamp and the previous modification timestamp have not changed, then no additional lines **20** have been added to the logfile **18-1** since the previous iteration, and the process returns to block **214** to wait for the next iteration.

(35) For purposes of illustration, assume that since the last iteration and prior to the time period beginning at time T2 the remote computing device **14-1** recorded, or logged, the additional lines **20-5-20-6** to the logfile **18-1**. This modification of the logfile **18-1** alters the modification timestamp in the inode data and thus, at block **216**, the computing device **12** determines that the modification timestamp in the inode data differs from the modification timestamp stored in the memory location **52**. The computing device **12** stores the new modification timestamp in the memory location **52**, and the memory location **52** now has the modification timestamp value **54-2** (block **218**). The computing device **12** copies the number of previous lines **20** stored in the memory

location **56** to a memory location **59** (block **220**). The memory location **59** has the value **60-1**.

(36) The computing device **12** determines the new current total number of lines **20-1-20-6** in the logfile **18-1** (block **222**), and stores the new current total number of lines **20-1-20-6** in the memory location **56** (block **224**). The memory location **56** has a value **58-2**. The computing device **12** determines the number of new lines **20-5-20-6** based on the current total number of lines **20-1-20-6** and the previous total number of lines **20-1-20-4**. In this example, the number of new lines **20-5-20-6** is two. The computing device **12** reads the number of new lines **20-5-20-6** from the remote computing device **14-1** and stores the new lines **20-5-20-6** in the memory **40** (block **226**). The computing device **12** sends the lines **20-5-20-6** to the centralized monitoring service **26** (block **228**). This completes the second iteration.

(37) The computing device **12** waits a predetermined amount of time (block **214**). After the predetermined amount of time lapses, at a time period beginning at time T3, the computing device **12** begins the process again and obtains the logfile meta data **22-1** associated with the logfile **18-1** (block **200**). The computing device **12** obtains the inode identifier from the inode data and compares the inode identifier to the previously stored inode identifier stored in the memory location **50** to determine if the logfile **18-1** is being accessed for the first time (block **202**). Assume for purposes of illustration that the current inode identifier matches the inode identifier stored in the memory location **50**. If the current inode identifier matches the inode identifier stored in the memory location **50**, the logfile **18-1** is not being accessed for the first time. The computing device **12** then extracts the most recent modification timestamp from the inode data and compares the most recent modification timestamp with the previous modification timestamp stored in the memory location **52** (block **216**). For purposes of illustration, assume that since the last iteration and prior to the time T3 the remote computing device **14-1** recorded, or logged, the additional lines **20-7-20-8** to the logfile **18-1**. This modification of the logfile **18-1** alters the modification timestamp in the inode data and thus, at block **216**, the computing device **12** determines that the modification timestamp in the inode data differs from the modification timestamp stored in the memory location **52**. The computing device **12** stores the new modification timestamp in the memory location **52**, and the memory location **52** now has the modification timestamp value **54-3** (block **218**). The computing device **12** copies the number of previous lines **20** stored in the memory location **56** to a memory location **59** (block **220**). The memory location **59** has the value **60-2**.

(38) The computing device **12** determines the new current total number of lines **20-1-20-8** in the logfile **18-1** (block **222**), and stores the new current total number of lines in the memory location **56** (block **224**). The memory location **56** has a value **58-3**. The computing device **12** determines the number of new lines **20-7-20-8** based on the current total number of lines **20-1-20-8** and the previous total number of lines **20-1-20-6**. In this example, the number of new lines **20-7-20-8** is two. The computing device **12** reads the number of new lines **20-7-20-8** from the remote computing device **14-1** and stores the new lines **20-7-20-8** in the memory **40** (block **226**). The computing device **12** sends the lines **20-7-20-8** to the centralized monitoring service **26** (block **228**). This completes the third iteration.

(39) This process may continue indefinitely, and may be performed for the remote computing device **14-2-14-2** consecutively, or in parallel with the remote computing device **14-1**. In this manner, the computing device **12** securely consolidates the logfile contents of the logfiles **18-1-18-3** on a continuous basis from the remote computing devices **14-1-14-3** to the centralized monitoring service **26**.

(40) While for purposes of illustration the process has been described in terms of a number of lines of a logfile **18**, the examples are not limited to using a number of lines of the logfile **18**, and can utilize any mechanism by which additional content that has been added to the logfile **18** since the last iteration can be determined. For example, in some examples the computing device **12** may utilize a size of the logfile **18**, and copy all logfile content that has been added to a logfile **18** since the previous iteration based on an increase in size of the logfile **18**. For example, at each iteration

the computing device **12** may determine and store a total size of the logfile **18** in the memory **40**. At the successive iteration, the computing device **12** may retrieve the previous size of the logfile **18** from the memory **40**, determine a current size of the logfile **18**, determine a new amount of data based on the previous size and the current size, and retrieve a most recent amount of data added to the logfile **18** equal to the new amount of data.

(41) For example, during one iteration a logfile **18** may be 434 KB in size. At a next iteration, the logfile **18** may have grown to 437 KB in size, and the computing device **12** may read the last 3 kB of data from the logfile **18** and send the data to the centralized monitoring service **26**.

(42) FIG. 5 is a block diagram of the computing device **12** suitable for implementing examples according to one example. The computing device **12** may comprise any computing or electronic device capable of including firmware, hardware, and/or executing software instructions to implement the functionality described herein, such as a computer server device, a desktop computing device, or the like. The computing device **12** includes a processor device **62**, the system memory **40**, and a system bus **64**. The system bus **64** provides an interface for system components including, but not limited to, the system memory **40** and the processor device **62**. The processor device **62** can be any commercially available or proprietary processor.

(43) The system bus **64** may be any of several types of bus structures that may further interconnect to a memory bus (with or without a memory controller), a peripheral bus, and/or a local bus using any of a variety of commercially available bus architectures. The system memory **40** may include non-volatile memory **66** (e.g., read-only memory (ROM), erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), etc.), and/or volatile memory **68** (e.g., random-access memory (RAM)). A basic input/output system (BIOS) **70** may be stored in the non-volatile memory **66** and can include the basic routines that help to transfer information between elements within the computing device **12**. The volatile memory **68** may also include a high-speed RAM, such as static RAM, for caching data.

(44) The computing device **12** may further include or be coupled to a non-transitory computer-readable storage medium or storage device **72**, which may comprise, for example, an internal or external hard disk drive (HDD) (e.g., enhanced integrated drive electronics (EIDE) or serial advanced technology attachment (SATA)), HDD (e.g., EIDE or SATA) for storage, flash memory, or the like. The storage device **72** and other drives associated with computer-readable media and computer-usable media may provide non-volatile storage of data, data structures, computer-executable instructions, and the like, including, for example, the configuration information **41**. Although the description of computer-readable media above refers to an HDD, it should be appreciated that other types of media that are readable by a computer, such as Zip disks, magnetic cassettes, flash memory cards, cartridges, and the like, may also be used in the operating environment, and, further, that any such media may contain computer-executable instructions for performing novel methods of the disclosed examples.

(45) A number of modules can be stored in the storage device **72** and in the volatile memory **68**, including an operating system **74** and one or more program modules **76**, which may implement the functionality described herein in whole or in part, including, for example, the functionality described herein with regard to the agent module **38**. It is to be appreciated that the examples can be implemented with various commercially available operating systems **74**.

(46) All or a portion of the examples may be implemented as a computer program product stored on a transitory or non-transitory computer-usable or computer-readable storage medium, such as the storage device **72**, which includes complex programming instructions, such as complex computer-readable program code, configured to cause the processor device **62** to carry out the steps described herein. Thus, the computer-readable program code can comprise software instructions for implementing the functionality of the examples described herein when executed on the processor device **62**. The processor device **62**, in conjunction with the program modules **76** in the volatile memory **68**, may serve as a controller, or control system, for the computing device **12** that is

configured to, or adapted to, implement the functionality described herein.

(47) An operator or other user may also be able to enter one or more configuration commands through a keyboard (not illustrated), a pointing device such as a mouse (not illustrated), or a touch-sensitive surface (not illustrated). Such input devices may be connected to the processor device **62** through an input device interface **78** that is coupled to the system bus **64** but can be connected by other interfaces such as a parallel port, an Institute of Electrical and Electronic Engineers (IEEE) 1394 serial port, a Universal Serial Bus (USB) port, an IR interface, and the like.

(48) The computing device **12** also includes a communication interface **80** suitable for communicating with the network **16** as appropriate or desired. The computing device **12** may also include a video port **82** configured to interface with a display **84**, to provide the operator or user information during the examples disclosed herein.

(49) FIG. **6** is a block diagram of a simplified version of the system **10** illustrated in FIG. **1**, according to one example. The system **10** includes the computing device **12** and the plurality of remote computing devices **14**. The computing device **12** includes the memory **40** and the processor device **62** coupled to the memory **40**. The processor device **62** establishes connections with the plurality of remote computing devices **14**. Each remote computing device **14** has a corresponding logfile **18**. For a plurality of iterations, the processor device **62** retrieves logfile contents from each logfile **18** on each remote computing device **14**, and sends the logfile contents to the centralized monitoring service **26**.

(50) Individuals will recognize improvements and modifications to the examples of the disclosure. All such improvements and modifications are considered within the scope of the concepts disclosed herein and the claims that follow.

Claims

1. A method for consolidating logfile contents from remote computing devices, comprising: establishing, by a computing device comprising a processor device via a network, connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile comprising information that the remote computing device is configured to store in the corresponding logfile; and for a plurality of iterations, by the computing device via the network: for each remote computing device, accessing logfile meta data that identifies attributes of the corresponding logfile, the logfile meta data being separate from data stored in the corresponding logfile, wherein the logfile meta data comprises inode data indicative of whether the corresponding logfile was previously accessed; and for each remote computing device, based on the inode data indicative of whether the corresponding logfile was accessed, retrieving at least a portion of logfile contents from the corresponding logfile, wherein retrieving the at least the portion of the logfile contents comprises: determining whether a first inode value of the inode data that identifies an inode at a time T1 and that refers to the corresponding logfile is different from a second inode value of the inode data that identifies the inode at a time T2; based on the first inode value being different from the second inode value, determining that the logfile was previously accessed and that the logfile has an old logfile status; retrieving the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed; and sending the at least the portion of the logfile contents to a centralized monitoring service.

2. The method of claim 1 wherein retrieving the at least the portion of the logfile contents further comprises determining a total number of lines of the logfile and storing the total number of lines in a memory.

3. The method of claim 2 wherein retrieving the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed comprises: retrieving a previous total number of lines in the logfile from the memory; determining a current total number of lines in the logfile; determining a number of new lines in the logfile based on the previous total

number of lines in the logfile and the current total number of lines in the logfile; and retrieving a number of most recent lines written to the logfile equal to the number of new lines.

4. The method of claim 1 further comprising: storing, for each logfile, the first inode value that identifies the inode at the time T1 that refers to the corresponding logfile.

5. The method of claim 1 wherein retrieving the at least the portion of the logfile contents further comprises determining a total size of the logfile and storing the total size of the logfile in a memory.

6. The method of claim 5 wherein retrieving the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed comprises: retrieving a previous size of the logfile from the memory; determining a current size of the logfile; determining a new amount of data based on the previous size and the current size; and retrieving a most recent amount of data added to the logfile equal to the new amount of data.

7. The method of claim 1 further comprising waiting a predetermined amount of time, and after the predetermined amount of time, performing another iteration of the plurality of iterations.

8. The method of claim 1 wherein establishing the connections with the plurality of remote computing devices comprises establishing secure connections with the plurality of remote computing devices.

9. The method of claim 8 wherein establishing secure connections with the plurality of remote computing devices comprises obtaining, for each remote computing device of the plurality of remote computing devices, a public key associated with the remote computing device, and establishing the secure connections with the plurality of remote computing devices using the public key associated with each respective remote computing device.

10. A computing device comprising: a memory; and a processor device coupled to the memory to: establish connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile comprising information that the remote computing device is configured to store in the corresponding logfile; and for a plurality of iterations: for each remote computing device, access logfile meta data that identifies attributes of the corresponding logfile, the logfile meta data being separate from data stored in the corresponding logfile, wherein the logfile meta data comprises inode data indicative of whether the corresponding logfile was previously accessed; and for each remote computing device, based on the inode data indicative of whether the corresponding logfile was accessed, retrieve at least a portion of logfile contents from the corresponding logfile, and wherein, to retrieve the at least the portion of the logfile contents, the processor device is further to: determine whether a first inode value of the inode data that identifies an inode at a time T1 and that refers to the corresponding logfile is different from a second inode value of the inode data that identifies the inode at a time T2; based on the first inode value being different from the second inode value, determine that the logfile was previously accessed and that the logfile has an old logfile status; retrieve the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed; and send the at least the portion of the logfile contents to a centralized monitoring service.

11. The computing device of claim 10 wherein, to retrieve the at least the portion of the logfile contents, the processor device is further to determine a total number of lines of the logfile, and store the total number of lines in a memory.

12. The computing device of claim 11 wherein to retrieve the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed the processor device is further to: retrieve a previous total number of lines in the logfile from the memory; determine a current total number of lines in the logfile; determine a number of new lines in the logfile based on the previous total number of lines in the logfile and the current total number of lines in the logfile; and retrieve a number of most recent lines written to the logfile equal to the number of new lines.

13. A computer program product for consolidating log information from remote computing devices,

the computer program product stored on a non-transitory computer-readable storage medium and including instructions to cause a processor device to carry out steps of: establishing connections with a plurality of remote computing devices, each remote computing device having a corresponding logfile comprising information that the remote computing device is configured to store in the corresponding logfile; for a plurality of iterations, by a computing device via a network: for each remote computing device, accessing logfile meta data that identifies attributes of the corresponding logfile, the logfile meta data being separate from data stored in the corresponding logfile, wherein the logfile meta data comprises inode data indicative of whether the corresponding logfile was previously accessed; and for each remote computing device, based on the inode data indicative of whether the corresponding logfile was accessed, retrieving at least a portion of logfile contents from the corresponding logfile, and wherein retrieving the at least the portion of the logfile contents comprises: determining whether a first inode value of the inode data that identifies an inode at a time T1 and that refers to the corresponding logfile is different from a second inode value of the inode data that identifies the inode at a time T2; based on the first inode value being different from the second inode value, determining that the logfile was previously accessed and that the logfile has an old logfile status; retrieving the at least the portion of the logfile contents of the logfile that have been added to the logfile since the logfile was previously accessed; and sending the at least the portion of the logfile contents to a centralized monitoring service.
