

# US Patent & Trademark Office

## Patent Public Search | Text View

---

|                      |                  |
|----------------------|------------------|
| United States Patent | 12393596         |
| Kind Code            | B2               |
| Date of Patent       | August 19, 2025  |
| Inventor(s)          | Cooley; Nicholas |

---

### Automated sampling of query results for training of a query engine

---

#### Abstract

An online system may generate numerous search records in response to searches requested by users. The online system may use a specific way to sample the historical search records to reduce biases in sampling. For example, the online system retrieves historical query records associated with an item query engine. The set of historical query records includes a plurality of search phrases. A historical query record is associated with a search phrase and a list of items returned by the item query engine. The online system determines the search frequencies for the search phrases. The online system stratifies the historical query records into a plurality of bins according to the search frequencies of the search phrases. The online system samples the historical query records from the plurality of bins to collect a representative set of historical query records and outputs the representative set of historical query records for rating.

---

|                   |   |
|-------------------|---|
| <b>Inventors:</b> | <b>Cooley; Nicholas (Oakland, CA)</b>     |
| <b>Applicant:</b> | <b>Maplebear Inc. (San Francisco, CA)</b> |
| <b>Family ID:</b> | <b>1000008763330</b>                      |
| <b>Assignee:</b>  | <b>Maplebear Inc. (San Francisco, CA)</b> |
| <b>Appl. No.:</b> | <b>18/590055</b>                          |
| <b>Filed:</b>     | <b>February 28, 2024</b>                  |

#### Prior Publication Data

|                            |                         |
|----------------------------|-------------------------|
| <b>Document Identifier</b> | <b>Publication Date</b> |
| US 20240202204 A1          | Jun. 20, 2024           |

#### Related U.S. Application Data

continuation parent-doc US 17826162 20220527 US 11947551 child-doc US 18590055

---

## Publication Classification

**Int. Cl.:** G06F16/24 (20190101); G06F16/2453 (20190101); G06F16/2455 (20190101); G06F16/2457 (20190101)

**U.S. Cl.:**

**CPC** G06F16/24578 (20190101); G06F16/24542 (20190101); G06F16/2455 (20190101);

## Field of Classification Search

**CPC:** G06F (16/24578); G06F (16/24542); G06F (16/2455); G06F (16/951); H04L (47/70)

**USPC:** 707/723

---

## References Cited

### U.S. PATENT DOCUMENTS

| Patent No.   | Issued Date | Patentee Name    | U.S. Cl.   | CPC          |
|--------------|-------------|------------------|------------|--------------|
| 5966126      | 12/1998     | Szabo            | N/A        | N/A          |
| 6418448      | 12/2001     | Sarkar           | N/A        | G06F 16/972  |
| 6601061      | 12/2002     | Holt             | N/A        | G06F 16/951  |
| 6952755      | 12/2004     | Karube           | 711/147    | G06F 9/52    |
| 7448022      | 12/2007     | Ram              | 717/120    | G06F 9/45504 |
| 7602834      | 12/2008     | Giallorenzi      | 375/150    | H04B 1/70752 |
| 7725496      | 12/2009     | Paval            | 707/791    | G06F 16/164  |
| 8046797      | 12/2010     | Bentolila et al. | N/A        | N/A          |
| 8180789      | 12/2011     | Wasserman        | 707/766    | G06F 16/242  |
| 8521526      | 12/2012     | Lloyd et al.     | N/A        | N/A          |
| 8554769      | 12/2012     | Thakur           | 707/735    | G06F 16/951  |
| 8751283      | 12/2013     | Bobak            | 705/348    | G06Q 10/06   |
| 9009146      | 12/2014     | Lopatenko et al. | N/A        | N/A          |
| 9208233      | 12/2014     | Tirumalareddy    | N/A        | G06F 16/951  |
| 9378065      | 12/2015     | Shear            | N/A        | G06Q 10/0631 |
| 9588648      | 12/2016     | Li               | N/A        | N/A          |
| 9779357      | 12/2016     | Dabell           | N/A        | N/A          |
| 9792160      | 12/2016     | Shear            | N/A        | G06F 9/50    |
| 9904579      | 12/2017     | Shear            | N/A        | G06F 16/245  |
| 10019486     | 12/2017     | Kephart et al.   | N/A        | N/A          |
| 10075384     | 12/2017     | Shear            | N/A        | H04L 63/10   |
| 10747793     | 12/2019     | Kataria et al.   | N/A        | N/A          |
| 10997244     | 12/2020     | Russell          | N/A        | G06F 16/27   |
| 11093556     | 12/2020     | Li               | N/A        | G06F 16/367  |
| 11507752     | 12/2021     | Bhowmik          | N/A        | G06F 40/295  |
| 11621899     | 12/2022     | Bettaiah et al.  | N/A        | N/A          |
| 2003/0208473 | 12/2002     | Lennon           | N/A        | G06F 16/71   |
| 2004/0064549 | 12/2003     | Brown            | 709/224    | H04L 43/00   |
| 2004/0095998 | 12/2003     | Luo              | 348/E5.066 | H04N 5/145   |
| 2005/0114756 | 12/2004     | Lehikoinen       | 715/208    | G06F 16/9562 |

|              |         |                   |          |               |
|--------------|---------|-------------------|----------|---------------|
| 2007/0255696 | 12/2006 | Desbarats         | N/A      | G06Q 30/02    |
| 2010/0202630 | 12/2009 | Lindahl           | 381/103  | H03G 5/165    |
| 2010/0318511 | 12/2009 | Phan              | 707/769  | G06F 16/9024  |
| 2011/0004621 | 12/2010 | Kelley            | 707/769  | G06F 16/21    |
| 2012/0317102 | 12/2011 | Hu et al.         | N/A      | N/A           |
| 2013/0290316 | 12/2012 | Bhatnagar et al.  | N/A      | N/A           |
| 2014/0026143 | 12/2013 | Suzuki            | 718/104  | G06F 9/52     |
| 2014/0052492 | 12/2013 | Boss              | 705/7.25 | G06Q 10/06375 |
| 2014/0074608 | 12/2013 | Beck et al.       | N/A      | N/A           |
| 2014/0250115 | 12/2013 | Yang              | N/A      | N/A           |
| 2014/0280193 | 12/2013 | Cronin et al.     | N/A      | N/A           |
| 2014/0280952 | 12/2013 | Shear             | 709/226  | H04L 63/10    |
| 2015/0033305 | 12/2014 | Shear             | 726/11   | G06F 21/53    |
| 2015/0163099 | 12/2014 | Huang             | 370/242  | H04L 12/40013 |
| 2015/0310115 | 12/2014 | Ryger             | 707/708  | G06F 16/9535  |
| 2016/0019338 | 12/2015 | Chudova           | 702/20   | G16B 30/10    |
| 2016/0034305 | 12/2015 | Shear             | 707/722  | G06F 16/285   |
| 2016/0217509 | 12/2015 | Eggleston et al.  | N/A      | N/A           |
| 2016/0266939 | 12/2015 | Shear             | N/A      | G06F 16/24575 |
| 2016/0299982 | 12/2015 | Bhave             | N/A      | G06F 16/3337  |
| 2017/0180910 | 12/2016 | Foti              | N/A      | H04W 4/70     |
| 2017/0193579 | 12/2016 | Goldberg          | N/A      | G06Q 30/0201  |
| 2017/0235816 | 12/2016 | Livshits et al.   | N/A      | N/A           |
| 2017/0242893 | 12/2016 | Kephart et al.    | N/A      | N/A           |
| 2018/0024901 | 12/2017 | Tankersley et al. | N/A      | N/A           |
| 2018/0089281 | 12/2017 | Li                | N/A      | G06F 16/9538  |
| 2018/0232461 | 12/2017 | Fu                | N/A      | N/A           |
| 2018/0278710 | 12/2017 | Sun               | N/A      | H04L 67/565   |
| 2018/0315145 | 12/2017 | Darnell et al.    | N/A      | N/A           |
| 2018/0357299 | 12/2017 | Miranda et al.    | N/A      | N/A           |
| 2019/0018904 | 12/2018 | Russell           | N/A      | G06F 16/9024  |
| 2019/0130050 | 12/2018 | Haas et al.       | N/A      | N/A           |
| 2019/0143999 | 12/2018 | Tsukidate         | 701/29.2 | B60R 16/02    |
| 2020/0192903 | 12/2019 | Bansal et al.     | N/A      | N/A           |
| 2020/0233893 | 12/2019 | He et al.         | N/A      | N/A           |
| 2020/0250197 | 12/2019 | Yang et al.       | N/A      | N/A           |
| 2020/0334263 | 12/2019 | Bansal et al.     | N/A      | N/A           |
| 2021/0081848 | 12/2020 | Polleri et al.    | N/A      | N/A           |
| 2021/0089531 | 12/2020 | Hu et al.         | N/A      | N/A           |
| 2021/0256068 | 12/2020 | Larlus            | N/A      | G06F 16/90332 |
| 2022/0050818 | 12/2021 | Li et al.         | N/A      | N/A           |
| 2022/0121884 | 12/2021 | Zadeh             | N/A      | G06N 3/006    |
| 2022/0122721 | 12/2021 | Alberico          | N/A      | N/A           |
| 2023/0385911 | 12/2022 | Weintraub         | N/A      | G06Q 30/0633  |

## OTHER PUBLICATIONS

PCT International Search Report and Written Opinion, PCT Application No. PCT/US22/49822, Feb. 22, 2023, 8 pages. cited by applicant

United States Office Action, U.S. Appl. No. 17/826,162, filed Jun. 8, 2023, 21 pages. cited by applicant

## Background/Summary

CROSS REFERENCE TO RELATED APPLICATION (1) This application is a continuation of U.S. Application Ser. No. 17/826,162, filed May 27, 2022, which is incorporated by reference in its entirety.

### FIELD

(1) This disclosure relates generally to training of a query search engine for a database and, more specifically, to sampling historical query records to assess training quality.

### BACKGROUND

(2) Delivering accurate, relevant and sometimes user-tailored data to users is often a challenging task for any online system, particularly ones with large databases. A large-scale database that serves millions of users often includes numerous data records and items and relies heavily on a query engine to return relevant and ranked results to the users. The performance of the query engine could significantly affect the user experience of an online system. The data retrieval task is even more complex in a database where the data is relatively dynamic. For example, in an inventory offering or management online system, the item availability and timing factors could affect the returned result generated by an item query engine. Additionally, result quality can vary subjectively from user to user. Since the query results in an item query engine could change based on various conditions of the inventory, evaluating the performance of such an item query engine is uniquely challenging.

### SUMMARY

(3) In some embodiments, a process of sampling historical query records for an item query engine is disclosed. The process may reduce bias in the sampling. In a large-scale online system, thousands or even millions of searches are performed by the online system. It is often infeasible to use all of the query records to evaluate the performance of the query engine. As such, an online system may sample the query records to generate a representative set. In an item query engine for an online system such as an inventory offering or management system, user searches could be skewed heavily on certain common items, such as common daily products. Hence, certain search phrases may predominate among the historical query records. A purely random sampling of the historical query records may generate bias in favor of the common search phrases because the randomly sampled collection will include a majority of records that are associated with those common search phrases.

(4) In some embodiments, an online system may retrieve a set of historical query records associated with an item query engine. The set of historical query records may include a plurality of search phrases. A historical query record is associated with a search phrase and a list of items returned by the item query engine. The online system may determine the search frequencies of the search phrases among the historical query records. The online system may stratify the set of historical query records into a plurality of bins according to the search frequencies of the search phrases. A bin corresponds to a subset of historical query records whose search phrases' search frequencies are within a range. The online system may sample the historical query records from the plurality of bins to collect a representative set of historical query records. The online system may output the representative set of historical query records for rating. The rated historical query records may be

used to evaluate the performance of the engine, conduct additional training, and refine the item query engine.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram of a system environment in which an online system, such as an online concierge system, operates, according to one or more embodiments.
- (2) FIG. 2 illustrates an environment of an online shopping concierge service, according to one or more embodiments.
- (3) FIG. 3 is a diagram of an online shopping concierge system, according to one or more embodiments.
- (4) FIG. 4A is a diagram of a customer mobile application (CMA), according to one or more embodiments.
- (5) FIG. 4B is a diagram of a shopper mobile application (SMA), according to one or more embodiments.
- (6) FIG. 5 is a block diagram of an item query engine, according to some embodiments.
- (7) FIG. 6 is a block diagram illustrating an example process of generating user item query results and training of the item query engine, according to one or more embodiments.
- (8) FIG. 7 is a conceptual diagram illustrating an example query log that includes a number of historical query records, according to one or more embodiments.
- (9) FIG. 8 is a flowchart depicting an example process for sampling historical query records, according to one or more embodiments.
- (10) FIG. 9 is a graphical illustration of stratifying the historical query records into different bins, according to one or more embodiments.
- (11) The figures depict embodiments of the present disclosure for purposes of illustration only. Alternative embodiments of the structures and methods illustrated herein may be employed without departing from the principles, or benefits touted, of the disclosure described herein.

### DETAILED DESCRIPTION

#### (12) System Overview

(13) FIG. 1 is a block diagram of a system environment **100** in which an online system, such as an online concierge system **102** as further described below in conjunction with FIGS. 2 and 3, operates. The system environment **100** shown by FIG. 1 comprises one or more client devices **110**, a network **120**, one or more third-party systems **130**, and the online concierge system **102**. In alternative configurations, different and/or additional components may be included in the system environment **100**. Additionally, in other embodiments, the online concierge system **102** may be replaced by an online system configured to retrieve content for display to users and to transmit the content to one or more client devices **110** for display.

(14) The client devices **110** are one or more computing devices capable of receiving user input as well as transmitting and/or receiving data via the network **120**. In one or more embodiments, a client device **110** is a computer system, such as a desktop or a laptop computer. Alternatively, a client device **110** may be a device having computer functionality, such as a personal digital assistant (PDA), a mobile telephone, a smartphone, or another suitable device. A client device **110** is configured to communicate via the network **120**. In one or more embodiments, a client device **110** executes an application allowing a user of the client device **110** to interact with the online concierge system **102**. For example, the client device **110** executes a customer mobile application **206** or a shopper mobile application **212**, as further described below in conjunction with FIGS. 4A and 4B, respectively, to enable interaction between the client device **110** and the online concierge system **102**. As another example, a client device **110** executes a browser application to enable

interaction between the client device **110** and the online concierge system **102** via the network **120**. In another embodiment, a client device **110** interacts with the online concierge system **102** through an application programming interface (API) running on a native operating system of the client device **110**, such as IOS® or ANDROID™

(15) A client device **110** includes one or more processors **112** configured to control operation of the client device **110** by performing functions. In various embodiments, a client device **110** includes a memory **114** comprising a non-transitory storage medium on which instructions are encoded. The memory **114** may have instructions encoded thereon that, when executed by the processor **112**, cause the processor to perform functions to execute the customer mobile application **206** or the shopper mobile application **212** to provide the functions further described above in conjunction with FIGS. **4A** and **4B**, respectively.

(16) The client devices **110** are configured to communicate via the network **120**, which may comprise any combination of local area and/or wide area networks, using both wired and/or wireless communication systems. In one or more embodiments, the network **120** uses standard communications technologies and/or protocols. For example, the network **120** includes communication links using technologies such as Ethernet, 802.11, worldwide interoperability for microwave access (WiMAX), 3G, 4G, 5G, code division multiple access (CDMA), digital subscriber line (DSL), etc. Examples of networking protocols used for communicating via the network **120** include multiprotocol label switching (MPLS), transmission control protocol/Internet protocol (TCP/IP), hypertext transport protocol (HTTP), simple mail transfer protocol (SMTP), and file transfer protocol (FTP). Data exchanged over the network **120** may be represented using any suitable format, such as hypertext markup language (HTML) or extensible markup language (XML). In some embodiments, all or some of the communication links of the network **120** may be encrypted using any suitable technique or techniques.

(17) One or more third party systems **130** may be coupled to the network **120** for communicating with the online concierge system **102** or with the one or more client devices **110**. In one or more embodiments, a third party system **130** is an application provider communicating information describing applications for execution by a client device **110** or communicating data to client devices **110** for use by an application executing on the client device. In other embodiments, a third party system **130** provides content or other information for presentation via a client device **110**. For example, the third party system **130** stores one or more web pages and transmits the web pages to a client device **110** or to the online concierge system **102**. The third party system **130** may also communicate information to the online concierge system **102**, such as advertisements, content, or information about an application provided by the third party system **130**.

(18) The online concierge system **102** includes one or more processors **142** configured to control operation of the online concierge system **102** by performing functions. In various embodiments, the online concierge system **102** includes a memory **144** comprising a non-transitory storage medium on which instructions are encoded. The memory **144** may have instructions encoded thereon corresponding to the modules further below in conjunction with FIG. **3** that, when executed by the processor **142**, cause the processor to perform the functionality and various processes further described in this disclosure, for example in conjunction with FIGS. **2**, **6** and **8**. For example, the memory **144** has instructions encoded thereon that, when executed by the processor **142**, cause the processor **142** to stratify historical query records of an item query engine according to the search frequencies of the search phrases in the queries and sample the historical query records accordingly to generate an unbiased representative set of historical query records for training of the item query engine. Additionally, the online concierge system **102** includes a communication interface configured to connect the online concierge system **102** to one or more networks, such as network **120**, or to otherwise communicate with devices (e.g., client devices **110**) connected to the one or more networks.

(19) One or more of a client device, a third party system **130**, or the online concierge system **102**

may be special purpose computing devices configured to perform specific functions, as further described below in conjunction with FIGS. 2-9, and may include specific computing components such as processors, memories, communication interfaces, and/or the like.

(20) System Overview

(21) FIG. 2 illustrates an environment **200** of an online platform, such as an online concierge system **102**, according to one or more embodiments. The figures use like reference numerals to identify like elements. A letter after a reference numeral, such as “**210a**,” indicates that the text refers specifically to the element having that particular reference numeral. A reference numeral in the text without a following letter, such as “**210**,” refers to any or all of the elements in the figures bearing that reference numeral. For example, “**210**” in the text refers to reference numerals “**210a**” or “**210b**” in the figures.

(22) The environment **200** includes an online concierge system **102**. The online concierge system **102** is configured to receive orders from one or more users **204** (only one is shown for the sake of simplicity). An order specifies a list of goods (items or products) to be delivered to the user **204**. The order also specifies the location to which the goods are to be delivered, and a time window during which the goods should be delivered. In some embodiments, the order specifies one or more retailers from which the selected items should be purchased. The user may use a customer mobile application (CMA) **206** to place the order; the CMA **206** is configured to communicate with the online concierge system **102**.

(23) The online concierge system **102** is configured to transmit orders received from users **204** to one or more shoppers **208**. A shopper **208** may be a contractor, employee, other person (or entity), robot, or other autonomous device enabled to fulfill orders received by the online concierge system **102**. The shopper **208** travels between a warehouse and a delivery location (e.g., the user's home or office). A shopper **208** may travel by car, truck, bicycle, scooter, foot, or other mode of transportation. In some embodiments, the delivery may be partially or fully automated, e.g., using a self-driving car. The environment **200** also includes three warehouses **210a**, **210b**, and **210c** (only three are shown for the sake of simplicity; the environment could include hundreds of warehouses). The warehouses **210** may be physical retailers, such as grocery stores, discount stores, department stores, etc., or non-public warehouses storing items that can be collected and delivered to users. Each shopper **208** fulfills an order received from the online concierge system **102** at one or more warehouses **210**, delivers the order to the user **204**, or performs both fulfillment and delivery. In one or more embodiments, shoppers **208** make use of a shopper mobile application **212** which is configured to interact with the online concierge system **102**.

(24) FIG. 3 is a diagram of an online concierge system **102**, according to one or more embodiments. In various embodiments, the online concierge system **102** may include different or additional modules than those described in conjunction with FIG. 3. Further, in some embodiments, the online concierge system **102** includes fewer modules than those described in conjunction with FIG. 3.

(25) The online concierge system **102** includes an inventory management engine **302**, which interacts with inventory systems associated with each warehouse **210**. In one or more embodiments, the inventory management engine **302** requests and receives inventory information maintained by the warehouse **210**. The inventory of each warehouse **210** is unique and may change over time. The inventory management engine **302** monitors changes in inventory for each participating warehouse **210**. The inventory management engine **302** is also configured to store inventory records in an inventory database **304**. The inventory database **304** may store information in separate records—one for each participating warehouse **210**—or may consolidate or combine inventory information into a unified record. Inventory information includes attributes of items that include both qualitative and qualitative information about items, including size, color, weight, SKU, serial number, and so on. In one or more embodiments, the inventory database **304** also stores purchasing rules associated with each item, if they exist. For example, age-restricted items

such as alcohol and tobacco are flagged accordingly in the inventory database **304**. Additional inventory information useful for predicting the availability of items may also be stored in the inventory database **304**. For example, for each item-warehouse combination (a particular item at a particular warehouse), the inventory database **304** may store a time that the item was last found, a time that the item was last not found (a shopper looked for the item but could not find it), the rate at which the item is found, and the popularity of the item.

(26) For each item, the inventory database **304** identifies one or more attributes of the item and corresponding values for each attribute of an item. For example, the inventory database **304** includes an entry for each item offered by a warehouse **210**, with an entry for an item including an item identifier that uniquely identifies the item. The entry includes different fields, with each field corresponding to an attribute of the item. A field of an entry includes a value for the attribute corresponding to the attribute for the field, allowing the inventory database **304** to maintain values of different categories for various items.

(27) In various embodiments, the inventory management engine **302** maintains a taxonomy of items offered for purchase by one or more warehouses **210**. For example, the inventory management engine **302** receives an item catalog from a warehouse **210** identifying items offered for purchase by the warehouse **210**. From the item catalog, the inventory management engine **302** determines a taxonomy of items offered by the warehouse **210**, different levels in the taxonomy providing different levels of specificity about items included in the levels. In various embodiments, the taxonomy identifies a category and associates one or more specific items with the category. For example, a category identifies “milk,” and the taxonomy associates identifiers of different milk items (e.g., milk offered by different brands, milk having one or more different attributes, etc.), with the category. Thus, the taxonomy maintains associations between a category and specific items offered by the warehouse **210** matching the category. In some embodiments, different levels in the taxonomy identify items with differing levels of specificity based on any suitable attribute or combination of attributes of the items. For example, different levels of the taxonomy specify different combinations of attributes for items, so items in lower levels of the hierarchical taxonomy have a greater number of attributes, corresponding to greater specificity in a category, while items in higher levels of the hierarchical taxonomy have a fewer number of attributes, corresponding to less specificity in a category. In various embodiments, higher levels in the taxonomy include less detail about items, so greater numbers of items are included in higher levels (e.g., higher levels include a greater number of items satisfying a broader category). Similarly, lower levels in the taxonomy include greater detail about items, so fewer numbers of items are included in the lower levels (e.g., higher levels include a fewer number of items satisfying a more specific category). The taxonomy may be received from a warehouse **210** in various embodiments. In other embodiments, the inventory management engine **302** applies a trained classification module to an item catalog received from a warehouse **210** to include different items in levels of the taxonomy, so application of the trained classification model associates specific items with categories corresponding to levels within the taxonomy.

(28) Inventory information provided by the inventory management engine **302** may supplement the training datasets **320**. Inventory information provided by the inventory management engine **302** may not necessarily include information about the outcome of picking a delivery order associated with the item, whereas the data within the training datasets **320** is structured to include an outcome of picking a delivery order (e.g., if the item in an order was picked or not picked).

(29) The online concierge system **102** also includes an order fulfillment engine **306** which is configured to synthesize and display an ordering interface to each user **204** (for example, via the customer mobile application **206**). The order fulfillment engine **306** is also configured to access the inventory database **304** in order to determine which products are available at which warehouse **210**. The order fulfillment engine **306** may supplement the product availability information from the inventory database **234** with an item availability predicted by the machine-learned item availability



model **316**. The order fulfillment engine **306** determines a sale price for each item ordered by a user **204**. Prices set by the order fulfillment engine **306** may or may not be identical to in-store prices determined by retailers (which is the price that users **204** and shoppers **208** would pay at the retail warehouses). The order fulfillment engine **306** also facilitates transactions associated with each order. In one or more embodiments, the order fulfillment engine **306** charges a payment instrument associated with a user **204** when he/she places an order. The order fulfillment engine **306** may transmit payment information to an external payment gateway or payment processor. The order fulfillment engine **306** stores payment and transactional information associated with each order in a transaction records database **308**.

(30) In various embodiments, the order fulfillment engine **306** generates and transmits a search interface to a client device of a user for display via the customer mobile application **106**. The order fulfillment engine **306** receives a query comprising one or more terms from a user and retrieves items satisfying the query, such as items having descriptive information matching at least a portion of the query. In various embodiments, the order fulfillment engine **306** leverages item embeddings for items to retrieve items based on a received query. For example, the order fulfillment engine **306** generates an embedding for a query and determines measures of similarity between the embedding for the query and item embeddings for various items included in the inventory database **304**.

(31) In some embodiments, the order fulfillment engine **306** also shares order details with warehouses **210**. For example, after successful fulfillment of an order, the order fulfillment engine **306** may transmit a summary of the order to the appropriate warehouses **210**. The summary may indicate the items purchased, the total value of the items, and in some cases, an identity of the shopper **208** and user **204** associated with the transaction. In one or more embodiments, the order fulfillment engine **306** pushes transaction and/or order details asynchronously to retailer systems. This may be accomplished via use of webhooks, which enable programmatic or system-driven transmission of information between web applications. In another embodiment, retailer systems may be configured to periodically poll the order fulfillment engine **306**, which provides detail of all orders which have been processed since the last request.

(32) The order fulfillment engine **306** may interact with a shopper management engine **310**, which manages communication with and utilization of shoppers **208**. In one or more embodiments, the shopper management engine **310** receives a new order from the order fulfillment engine **306**. The shopper management engine **310** identifies the appropriate warehouse **210** to fulfill the order based on one or more parameters, such as a probability of item availability determined by a machine-learned item availability model **316**, the contents of the order, the inventory of the warehouses, and the proximity to the delivery location. The shopper management engine **310** then identifies one or more appropriate shoppers **208** to fulfill the order based on one or more parameters, such as the shoppers' proximity to the appropriate warehouse **210** (and/or to the user **204**), his/her familiarity level with that particular warehouse **210**, and so on. Additionally, the shopper management engine **310** accesses a shopper database **312** which stores information describing each shopper **208**, such as his/her name, gender, rating, previous shopping history, and so on.

(33) As part of fulfilling an order, the order fulfillment engine **306** and/or shopper management engine **310** may access a user database **314** which stores information describing each user. This information could include each user's name, address, gender, shopping preferences, favorite items, stored payment instruments, and so on.

(34) In various embodiments, the order fulfillment engine **306** determines whether to delay display of a received order to shoppers for fulfillment by a time interval. In response to determining to delay the received order by a time interval, the order fulfillment engine **306** evaluates orders received after the received order and during the time interval for inclusion in one or more batches that also include the received order. After the time interval, the order fulfillment engine **306** displays the order to one or more shoppers via the shopper mobile application **212**; if the order fulfillment engine **306** generated one or more batches including the received order and one or more

orders received after the received order and during the time interval, the one or more batches are also displayed to one or more shoppers via the shopper mobile application **212**.

(35) The online concierge system **102** further includes a machine-learned item availability model **316**, a modeling engine **318**, and training datasets **320**. The modeling engine **318** uses the training datasets **320** to generate the machine-learned item availability model **316**. The machine-learned item availability model **316** can learn from the training datasets **320**, rather than follow only explicitly programmed instructions. The inventory management engine **302**, order fulfillment engine **306**, and/or shopper management engine **310** can use the machine-learned item availability model **316** to determine a probability that an item is available at a warehouse **210**. The machine-learned item availability model **316** may be used to predict item availability for items being displayed to or selected by a user or included in received delivery orders. A single machine-learned item availability model **316** is used to predict the availability of any number of items.

(36) The machine-learned item availability model **316** can be configured to receive as inputs information about an item, the warehouse for picking the item, and the time for picking the item. The machine-learned item availability model **316** may be adapted to receive any information that the modeling engine **318** identifies as indicators of item availability. At minimum, the machine-learned item availability model **316** receives information about an item-warehouse pair, such as an item in a delivery order and a warehouse at which the order could be fulfilled. Items stored in the inventory database **304** may be identified by item identifiers. As described above, various characteristics, some of which are specific to the warehouse (e.g., a time that the item was last found in the warehouse, a time that the item was last not found in the warehouse, the rate at which the item is found, the popularity of the item) may be stored for each item in the inventory database **304**. Similarly, each warehouse may be identified by a warehouse identifier and stored in a warehouse database along with information about the warehouse. A particular item at a particular warehouse may be identified using an item identifier and a warehouse identifier. In other embodiments, the item identifier refers to a particular item at a particular warehouse, so that the same item at two different warehouses is associated with two different identifiers. For convenience, both of these options to identify an item at a warehouse are referred to herein as an “item-warehouse pair.” Based on the identifier(s), the online concierge system **102** can extract information about the item and/or warehouse from the inventory database **304** and/or warehouse database and provide this extracted information as inputs to the item availability model **316**.

(37) The machine-learned item availability model **316** contains a set of functions generated by the modeling engine **318** from the training datasets **320** that relate the item, warehouse, and timing information, and/or any other relevant inputs, to the probability that the item is available at a warehouse. Thus, for a given item-warehouse pair, the machine-learned item availability model **316** outputs a probability that the item is available at the warehouse. The machine-learned item availability model **316** constructs the relationship between the input item-warehouse pair, timing, and/or any other inputs and the availability probability (also referred to as “availability”) that is generic enough to apply to any number of different item-warehouse pairs. In some embodiments, the probability output by the machine-learned item availability model **316** includes a confidence score. The confidence score may be the error or uncertainty score of the output availability probability and may be calculated using any standard statistical error measurement. In some examples, the confidence score is based in part on whether the item-warehouse pair availability prediction was accurate for previous delivery orders (e.g., if the item was predicted to be available at the warehouse and not found by the shopper or predicted to be unavailable but found by the shopper). In some examples, the confidence score is based in part on the age of the data for the item, e.g., if availability information has been received within the past hour, or the past day. The set of functions of the item availability model **316** may be updated and adapted following retraining with new training datasets **320**. The machine-learned item availability model **316** may be any machine learning model, such as a neural network, boosted tree, gradient boosted tree or random

forest model. In some examples, the machine-learned item availability model **316** is generated from XGBoost algorithm.

(38) The item probability generated by the machine-learned item availability model **316** may be used to determine instructions delivered to the user **204** and/or shopper **208**, as described in further detail below.

(39) The training datasets **320** relate a variety of different factors to known item availabilities from the outcomes of previous delivery orders (e.g., if an item was previously found or previously unavailable). The training datasets **320** include the items included in previous delivery orders, whether the items in the previous delivery orders were picked, warehouses associated with the previous delivery orders, and a variety of characteristics associated with each of the items (which may be obtained from the inventory database **304**). Each piece of data in the training datasets **320** includes the outcome of a previous delivery order (e.g., if the item was picked or not). The item characteristics may be determined by the machine-learned item availability model **316** to be statistically significant factors predictive of the item's availability. For different items, the item characteristics that are predictors of availability may be different. For example, an item type factor might be the best predictor of availability for dairy items, whereas a time of day may be the best predictive factor of availability for vegetables. For each item, the machine-learned item availability model **316** may weight these factors differently, where the weights are a result of a “learning” or training process on the training datasets **320**. The training datasets **320** are very large datasets taken across a wide cross section of warehouses, shoppers, items, warehouses, delivery orders, times, and item characteristics. The training datasets **320** are large enough to provide a mapping from an item in an order to a probability that the item is available at a warehouse. In addition to previous delivery orders, the training datasets **320** may be supplemented by inventory information provided by the inventory management engine **302**. In some examples, the training datasets **320** are historic delivery order information used to train the machine-learned item availability model **316**, whereas the inventory information stored in the inventory database **304** include factors input into the machine-learned item availability model **316** to determine an item availability for an item in a newly received delivery order. In some examples, the modeling engine **318** may evaluate the training datasets **320** to compare a single item's availability across multiple warehouses to determine if an item is chronically unavailable. This may indicate that an item is no longer manufactured. The modeling engine **318** may query a warehouse **210** through the inventory management engine **302** for updated item information on these identified items.

(40) The training datasets **320** include a time associated with previous delivery orders. In some embodiments, the training datasets **320** include a time of day at which each previous delivery order was placed. Time of day may impact item availability, since during high-volume shopping times, items may become unavailable that are otherwise regularly stocked by warehouses. In addition, availability may be affected by restocking schedules, e.g., if a warehouse mainly restocks at night, item availability at the warehouse will tend to decrease over the course of the day. Additionally, or alternatively, the training datasets **320** include a day of the week previous delivery orders were placed. The day of the week may impact item availability since popular shopping days may have reduced inventory of items or restocking shipments may be received on particular days. In some embodiments, training datasets **320** include a time interval since an item was previously picked in a previous delivery order. If an item has recently been picked at a warehouse, this may increase the probability that it is still available. If there has been a long time interval since an item has been picked, this may indicate that the probability that it is available for subsequent orders is low or uncertain. In some embodiments, training datasets **320** include a time interval since an item was not found in a previous delivery order. If there has been a short time interval since an item was not found, this may indicate that there is a low probability that the item is available in subsequent delivery orders. And conversely, if there has been a long time interval since an item was not found, this may indicate that the item may have been restocked and is available for subsequent delivery

orders. In some examples, training datasets **320** may also include a rate at which an item is typically found by a shopper at a warehouse, a number of days since inventory information about the item was last received from the inventory management engine **302**, a number of times an item was not found in a previous week, or any number of additional rate or time information. The relationships between this time information and item availability are determined by the modeling engine **318** training a machine learning model with the training datasets **320**, producing the machine-learned item availability model **316**.

(41) The training datasets **320** include item characteristics. In some examples, the item characteristics include a department associated with the item. For example, if the item is yogurt, it is associated with the dairy department. The department may be the bakery, beverage, nonfood, and pharmacy, produce and floral, deli, prepared foods, meat, seafood, dairy, the meat department, or dairy department, or any other categorization of items used by the warehouse. The department associated with an item may affect item availability, since different departments have different item turnover rates and inventory levels. In some examples, the item characteristics include an aisle of the warehouse associated with the item. The aisle of the warehouse may affect item availability since different aisles of a warehouse may be more frequently re-stocked than others. Additionally, or alternatively, the item characteristics include an item popularity score. The item popularity score for an item may be proportional to the number of delivery orders received that include the item. An alternative or additional item popularity score may be provided by a retailer through the inventory management engine **302**. In some examples, the item characteristics include a product type associated with the item. For example, if the item is a particular brand of a product, then the product type will be a generic description of the product type, such as “milk” or “eggs.” The product type may affect the item availability, since certain product types may have a higher turnover and re-stocking rate than others or may have larger inventories in the warehouses. In some examples, the item characteristics may include a number of times a shopper was instructed to keep looking for the item after he or she was initially unable to find the item, a total number of delivery orders received for the item, whether or not the product is organic, vegan, gluten free, or any other characteristics associated with an item. The relationships between item characteristics and item availability are determined by the modeling engine **318** training a machine learning model with the training datasets **320**, producing the machine-learned item availability model **316**.

(42) The training datasets **320** may include additional item characteristics that affect the item availability and can therefore be used to build the machine-learned item availability model **316** relating the delivery order for an item to its predicted availability. The training datasets **320** may be periodically updated with recent previous delivery orders. The training datasets **320** may be updated with item availability information provided directly from shoppers **208**. Following updating of the training datasets **320**, a modeling engine **318** may retrain a model with the updated training datasets **320** and produce a new machine-learned item availability model **316**.

(43) The item query engine **330** receives search queries from users and selects items to be presented as search results to users. Items and products may be used interchangeably in this disclosure. The item query engine **330** uses one or more machine learning models that are trained to select, score, and rank items. The item query engine **330** may be applied to different warehouses with different item selections and availabilities. In one or more embodiments, the item query engine **330** receives a search phrase from a user that includes one or more keywords. The item query engine **330** selects one or more items that match the search phrase. The item query engine **330** consults machine-learned item availability model **316** to determine the availabilities of those items. For available items, the item query engine **330** ranks and scores the items based on different criteria, such as relevancy, diversity, and whether an item is sponsored. The item query engine **330** in turn produces the result to a graphical user interface for the user to select the items. The item query engine **330** is discussed in further detail below in the context of FIG. 5. U.S. patent application Ser. No. 17/550,950, entitled Context-Based Content-Scoring for an Online Concierge

System, filed on Dec. 14, 2021, is incorporated by reference herein in its entirety for all purposes.

(44) Query logs **340** are stored in a database that saves historical query records of the online concierge system **102** for various warehouses **210**. A query log **340** may include a number of historical query records. A historical query record may be a record that documents an actual search requested by a user, time of the search, identifier of the warehouse **210** and the actual search result returned by the item query engine **330**. Each historical query record may be associated with a unique search identifier and a timestamp.

#### (45) Customer Mobile Application

(46) FIG. 4A is a diagram of the customer mobile application (CMA) **206**, according to one or more embodiments. The CMA **206** includes an ordering interface **402**, which provides an interactive interface with which the user **104** can browse through and select products and place an order. The CMA **206** also includes a system communication interface **404** which, among other functions, receives inventory information from the online shopping concierge system **102** and transmits order information to the system **102**. The CMA **206** also includes a preferences management interface **406** which allows the user **104** to manage basic information associated with his/her account, such as his/her home address and payment instruments. The preferences management interface **406** may also allow the user to manage other details such as his/her favorite or preferred warehouses **210**, preferred delivery times, special instructions for delivery, and so on.

#### (47) Shopper Mobile Application

(48) FIG. 4B is a diagram of the shopper mobile application (SMA) **212**, according to one or more embodiments. The SMA **212** includes a barcode scanning module **420** which allows a shopper **208** to scan an item at a warehouse **210** (such as a can of soup on the shelf at a grocery store). The barcode scanning module **420** may also include an interface which allows the shopper **108** to manually enter information describing an item (such as its serial number, SKU, quantity and/or weight) if a barcode is not available to be scanned. SMA **212** also includes a basket manager **422** which maintains a running record of items collected by the shopper **208** for purchase at a warehouse **210**. This running record of items is commonly known as a “basket.” In one or more embodiments, the barcode scanning module **420** transmits information describing each item (such as its cost, quantity, weight, etc.) to the basket manager **422**, which updates its basket accordingly. The SMA **212** also includes a system communication interface **424** which interacts with the online shopping concierge system **102**. For example, the system communication interface **424** receives an order from the online concierge system **102** and transmits the contents of a basket of items to the online concierge system **102**. The SMA **212** also includes an image encoder **426** which encodes the contents of a basket into an image. For example, the image encoder **426** may encode a basket of goods (with an identification of each item) into a QR code which can then be scanned by an employee of the warehouse **210** at check-out.

#### Example Item Query Engine

(49) FIG. 5 is a block diagram for an item query engine **330**, according to some embodiments. In one or more embodiments, the item query engine **330** may include a user embedding engine **500**, a query embedding engine **510**, an item embedding engine **520**, an anchor embedding engine **530**, an item scoring engine **540**, and a training engine **550**. Alternative embodiments may include more, fewer, or different components from those illustrated in FIG. 5, and the functionality of each component may be divided between the components differently than described in the description below. Additionally, each component may perform their respective functionalities in response to a request from a human, or automatically without human intervention.

(50) The user embedding engine **500** generates user embeddings for users using the online concierge system **102**. Each user embedding is an embedding vector that describes characteristics or features about an associated user. The user embedding engine **500** generates user embeddings for a user based on user data. User data is data that describes characteristics about a user that may be relevant for determining the relevance of a product to a user, and such data may be collected in

accordance with one or more privacy policies and/or applicable privacy laws and/or regulations. For example, user data may include one or more of the user's name, the user's location, the user's stated preferences, the user's previously ordered products, the user's frequency of placing orders, which retailers the user orders from, a typical order cost for the user, or a browsing history of the user on the customer mobile application **106** or other applications the user may use. User data may include raw data, preprocessed data, or feature sets describing information about a user. In some embodiments, the user embedding engine **500** collects user data from the user database **314**. The user embedding engine **500** may store generated user embeddings in the user database **314** and associate, within the user database **314**, each user embedding with a user that the user embedding describes.

(51) In some embodiments, the user embedding engine **500** uses one or more user models to generate user embeddings. User models are machine learning models (e.g., neural networks) that are trained to generate user embeddings based on user data. The user embedding engine **500** can also retrieve a user embedding associated with a specified user from the user database **314**.

(52) The query embedding engine **510** generates a query embedding for the user's search query. A query embedding is an embedding vector that describes features of the user's search query. The query embedding engine generates a query embedding based on search query data. Query data is data describing a user's search query for the online concierge system **102**. For example, query data may include one or more of search phrases, previous searches by the user within the user's session, or search queries conducted by other users of the online concierge system **102**. Query data may also include context data describing the context in which the user has queried the online concierge system **102** for products. For example, the context data may include one or more of how long the user's session with the online concierge system **102** has lasted, the products that are currently in the user's selected products list, or other products with which the user has interacted during the session. Query data may include raw data, preprocessed data, or feature sets describing information about a query or context. Any and/or all of this data may be collected in accordance with one or more privacy policies and/or applicable privacy laws and/or regulations.

(53) The query embedding engine **510** uses one or more query embedding models to generate a query embedding. Query embedding models are machine learning models (e.g., neural networks) that are trained to generate query embeddings based on search phrases.

(54) The item embedding engine **520** generates an item embedding for items being evaluated by the item query engine **330**. An item embedding is an embedding vector that describes an item. The item embeddings may be associated with specific items stored by the inventory database **304**. For example, each brand of a product may have an individual item embedding, or products may have different item embeddings for each retailer that sells the product. Alternatively, each item embedding may be associated with a generic product, and each generic product may be associated with specific products that are similar or substitutes of each other. For example, the inventory database **304** may store an item embedding for the generic product "milk", and the specific products of "Moo Moo 3% Milk" and "Greener Pastures Organic Whole Milk" may both be associated with the item embedding for "milk." In some embodiments, item embeddings are stored in the inventory database **304**.

(55) The item embedding engine **520** uses one or more item embedding models to generate an item embedding based on product data. An item embedding model is one or more machine learning models (e.g., neural networks) that are trained to generate item embeddings based on product data. Item data is data that describes characteristics about items available for purchase using the online concierge system **102**. For example, item data may include one or more of a product name, a product type, whether a product is associated with a recipe, retailers that offer the product for sale, the shelf-life of a product, identifiers for other products with which the product is commonly purchased, a popularity of the product, the availability of a product, the price of a product, any restrictions that may be in place on the purchase of the product, whether the product is a food item,

a frequency with which the product is purchased using the online concierge system **102**, other products which the product has been or may be presented, or an expense incurred by the online concierge system **102** to provide the product to the user. Product data may include raw data, preprocessed data, or feature sets describing information about a product.

(56) The anchor embedding engine **530** generates anchor embeddings based on user embeddings and query embeddings. An anchor embedding is an embedding of the same dimension and in the same embedding space as an item embedding. The anchor embedding can therefore be compared to item embeddings to determine products that would be relevant to present to a user in response to a search query. The anchor embedding engine **530** may use one or more anchor embedding models to generate an anchor embedding. Anchor embedding models are machine learning models (e.g., neural networks) that are trained to generate anchor embeddings based on user embeddings and query embeddings.

(57) The item scoring engine **540** generates item scores for items. An item score is a score for a product that indicates the product's affinity for being presented to a user in response to a search query from the user. An item score may represent a likelihood that the user will interact with the product if the product is presented to the user or may represent some expected value based on the likelihood of user interaction and the value of the user's interaction with the product. The item scoring engine **540** generates item scores for products based on a comparison of an anchor embedding with a set of item embeddings associated with a set of candidate products. The set of candidate products can include all products available on the online concierge system **102** or a subset of the products. The anchor embedding is generated by the anchor embedding engine **530** based on a query embedding for the search query and a user embedding for the user who submitted the search query. The item scoring engine **540** may compare the anchor embedding with the set of item embeddings by calculating a Euclidean distance, a cosine distance, or a dot product of the anchor embedding and each item embedding.

(58) Additionally, the item scoring engine **540** may use a machine learning model (e.g., a neural network) trained to generate item scores for products based on item embeddings associated with the products and an anchor embedding. In some embodiments, the machine learning model generates item scores based on comparisons of item embeddings for the set of candidate products with an anchor embedding.

(59) The online concierge system **102** may generate various embedding vectors using various machine learning models and techniques. In some embodiments, the words of the textual content are mapped into vectors using different embedding techniques such as term frequency-inverse document frequency (TF-IDF) vectorization, continuous big-of-words (CBOW) model, and/or skip-gram model. The mapping process may be conducted through a supervised or unsupervised neural network. The generation of the word vectors is based on aggregated word-to-word co-occurrence statistics from a corpus. A corpus may be selected from a collection of open-source data sources, a collection of textual content of data specific to the online concierge system **102** and may additionally include other sources of text from books, publications, online articles, advertisements, etc. to provide additional training to a neural network that performs the word vector generation. Each word vector generated corresponds to a word and represents the semantic correlation, similarity, and difference of the word with respect to other words in the corpus. Techniques such as TF-IDF vectorization may be used to penalize the weight of common words such as articles, prepositions, and conjunctions that carry little significance in defining semantic characteristics of a text.

(60) The generation of an embedding of a collection of words or data may be achieved in different ways. In one or more embodiments, after single words are converted into word vectors, an average of all of the word vectors can be calculated to generate a common vector that has a specific direction and magnitude. The average can be a simple average or a weighted average. For example, the weighted average can be calculated based on the number of occurrences of a word in the survey

response. Techniques such as TF-IDF vectorization may be used to reduce the weight of common words that do not carry much semantic significance. The averaged vector represents an overall semantic characteristic of the textual content of the survey response in the form of a mathematical vector. Such an averaged vector is served as the embedding of the collection of words or data. In another case, instead of taking the average of all word vectors, a certain number (e.g. ten) of top semantically significant word vectors are selected to generate an averaged vector.

(61) In another embodiment, the generation of an embedding for a collection of words or data may be carried out through another neural network. The neural network can be a deep neural network that includes an input layer, an output layer, and one or more hidden intermediate layers. Each layer includes one or more nodes that are connected to other layers. A layer receives inputs from a preceding layer and produces outputs for a succeeding layer. In one case earlier layers (e.g., layers closer to the input layer) are configured to capture syntactic meanings of the textual content, while later layers (i.e. layers closer to the output layer) are configured to capture semantic meanings of the textual content. The layers of the neural network perform recognition of syntactic and/or semantic features by convolution, clustering, classification, matching, and/or the like. The neural network is configured to receive the textual content. The neural network is configured to output a vector that represents the semantic characteristic of the textual content of a collection of words or data after the input is analyzed through multiple layers and nodes. The output vector represents the semantic characteristic of the textual content and may be served as the output embedding.

(62) The training engine **550** evaluates the performance of the machine learning models used by the item query engine **330** so that the machine learning models can be trained to perform more effectively. For example, the training engine **512** may evaluate the performance of user embedding models, query embedding models, item embedding models, or anchor embedding models. Periodically, historical query records may also be sampled and rated manually or according to one or more metrics. The rated historical query records may also be used to further train one or more machine learning models in the item query engine **330**. The periodic sampling of query records and training of item query engine **330** are further discussed in FIG. **6** through FIG. **9**.

(63) In some embodiments, the training engine **550** uses one or more loss functions to train machine learning models used by the item query engine **330**. A loss function may improve the machine learning models used by the item query engine **330** such that the item query engine **330** generates item scores for products representing an affinity of the products being presented to a user in response to a search query from the user. In some embodiments, the training engine **550** backpropagates the loss function through the anchor embedding models, the item embedding models, the query embedding models, and the user embedding models. The training engine **550** also may apply a machine learning model to the anchor embedding and a set of item embeddings to train the machine learning models used by the item query engine **330**.

(64) In some embodiments, the training engine **550** uses a loss function that is based on the triplet loss function. The modified triplet loss function may assign different weights to different types of interactions that are considered positive examples based on a hierarchy of interactions. The hierarchy of interactions may reflect that certain interactions are more valuable to the online concierge system **102** than other interactions. For example, a purchase interaction (e.g., where a user purchases a product) may be more valuable than a select interaction (e.g., where the user adds the product to a selected products list), which itself may be more valuable than a click interaction (e.g., where the user selects a product to see more information about the product). The training engine **550** may use weights in a modified triplet loss function that emphasizes minimizing the distance from higher tier positive examples than from lower tier positive examples. For example, the training engine may use the following as a loss function:

(65)

$$L = \sum_{i=1}^M \max(\delta_{i_1} \beta_1 \text{dist}(\text{AE}, \text{PE}_{P_{i_1}}) + \text{.Math.} + \delta_{i_n} \beta_n \text{dist}(\text{AE}, \text{PE}_{P_{i_n}}) - \text{dist}(\text{AE}, \text{PE}_{N_{i_n}}) + \alpha, 0)$$



(66) For this example loss function,  $AE$  is an anchor embedding for a user who submitted a search query,  $PE_{sub.p}$  is an item embedding of a positive training example (e.g., where the user interacted with a presented product), and  $PE_{sub.N}$  is an item embedding of a negative training example (e.g., where the user did not interact with a presented product). The function  $dist$  is a function that computes the distance between two embeddings (e.g., the Euclidean distance). In the example loss function, there are  $n$  types of positive interactions that a user can take with regards to a product and those different types of interactions are weighted using weights  $\beta_{sub.1} \dots \beta_{sub.n}$ . These weights can assign greater value to more valuable interactions by increasing the loss value of a distance between the anchor embedding and the positive example. Additionally, the example loss function uses indicator variables  $\delta_{sub.1} \dots \delta_{sub.n}$  to apply the correct weight  $\beta$  to the positive example. The only indicator variable with a non-zero value is the indicator variable that corresponds to the type of interaction represented by the positive example. Furthermore,  $M$  is the number of search examples used to train the machine learning models and  $\alpha$  is a hyperparameter controlling the margin in the example loss function.

#### Example Item Query and Engine Training

(67) FIG. 6 is a block diagram illustrating an example overall process of generating item query results and training of the item query engine 330, according to one or more embodiments. In various embodiments, the process may include different or additional steps than those described in conjunction with FIG. 6. Further, in some embodiments, the steps of the process may be performed in different orders than the order described in conjunction with FIG. 6. The process described in conjunction with FIG. 6 may be carried out by the online concierge system 102, while in other embodiments, the steps of the method are performed by any online system capable of generating query results in response to a user query.

(68) In one or more embodiments, various end users may use the online concierge system 102 to search for various products in different warehouses 210. A user query 610 may be a search phrase that includes one or more keywords. A user may enter the search phrase via a user interface such as the ordering interface 402 in a CMA 206. The online concierge system 102 receives the search phrase and uses the item query engine 330 to generate a query result 620 for the user.

(69) In various embodiments, the item query engine 330 may generate a query result 620 based on various factors. An example of the item query engine 330 and how it operates is further described above in association with FIG. 5. In some embodiments, a query result may include a list of items that are ranked for the positions for display. The generation of a query result may include a search phase and a rank phase. In the search phase, a search may be conducted for a collection of items in warehouse 210 that are relevant to the search phase. In some embodiments, a warehouse 210 may maintain a real-time list of inventory items that are available for purchase. In such a case, the item query engine 330 may remove unavailable items. In some embodiments, the online concierge system 102 may rely on the machine-learned item availability model 316 to predict whether an item is available. In such a case, the availability of an item may be one of the factors that determines whether an item is selected and how it is ranked. Other factors may also be used to select and rank the items, such as the popularity of the items, historical preference of the user, historical purchase of the user, profitability of items, expiration dates, diversity of items, etc. In some embodiments, a selection factor may also be whether an item is sponsored (e.g., an advertised item). In some instances, an item may be sponsored by a grocery store or other retailer using the service of the online concierge system 102 and/or by suppliers and/or manufacturers (e.g., a consumer packaged good company) that want to promote the item.

(70) The item query engine 330 may also rank the selected items to determine how the items are arranged in the user interface to be displayed to the user. The ranking may also be based on the various factors described above. Items that are ranked higher are displayed at more prominent positions in the query results, such as at the top of the results. In some embodiments, one or more positions may be reserved for certain types of items. For example, in some embodiments, certain

positions may be reserved for sponsored items and may each be associated with a different bid rate for an item to be placed at such a reserved position.

(71) The query results **620** may be timestamped and each assigned with a unique search identifier. The query results **620** that have occurred at the online concierge system **102** may be referred to as historical query records. The historical query records may be stored in one or more query logs **340**. A historical query record is associated with a search phrase and a list of items returned by the item query engine **330**. In some embodiments, the list may be ranked and the ranking is also saved as part of the historical query record. The query logs **340** may be stored in a database for future retrieval. An example of a query log is discussed in FIG. 7.

(72) In some embodiments, the online concierge system **102** may automatically and periodically sample **630** a representative set of historical query records for the rating of the query results as a part of a process of training and evaluating the performance of the item query engine **330**. In some embodiments, the online concierge system **102** receives numerous search requests from different users in various warehouses **210**. As such, the online concierge system **102** uses one or more algorithms to sample some of the historical query records for further evaluation. The sampling process **630** described in this disclosure may reduce the bias of the samples and improve the performance of the item query engine **330**.

(73) For example, in one or more embodiments, the online concierge system **102** conducts the sampling **630** of the historical query records based on the search frequencies of the search phrases. If the historical query records are sampled randomly based on the search identifiers, oftentimes frequently searched phrases such as “milk,” “apple,” etc. may be overly represented in the sampled result because the historical query records in a query log **340** may be predominantly filled with historical query records that are associated with those search phrases. In some embodiments, the online concierge system **102** may determine the search frequencies of the search phrases among the historical query records. The online concierge system **102** may stratify the historical query records into a plurality of bins according to the search frequencies of the search phrases. For example, a bin corresponds to a subset of historical query records whose search phrases' search frequencies are within a range. Each bin may have a different range than another bin. In turn, the online concierge system **102** samples the historical query records from the plurality of bins to collect a representative set of historical query records. The sampling process **630** is further described in FIGS. 8 and 9.

(74) The online concierge system **102** outputs the representative set of historical query records for rating **640** of the historical query records. The rating of the historical query records may be conducted by the online concierge system **102** or by another system. For example, the online concierge system **102** may transmit representative sets of historical query records to a third party system (e.g., a rating agent) through an API. In some embodiments, the online concierge system **102** periodically samples (e.g., every week, every N days, etc.) the historical query records and collect the sampled results for rating.

(75) In various embodiments, the rating may be performed in various ways, such as automatically or manually (e.g., manually reviewed by a person who evaluates the quality of the query result). The online concierge system **102** may define one or more metrics whose values are used to represent the quality of the query results. Example metrics may include relevancy, diversity, user preference, conversion rate (e.g., user click rate, user purchase rate) and other statistically or not, formula-based or not, subjective or objective, manually assigned or automatically assigned, quantifiable or not, metrics. In some embodiments, some of the evaluations may be automatically determined by the online concierge system **102** (e.g., past conversion rates of the item in a query result) while other evaluations may be conducted by a third-party rater that performs manual rating. In some embodiment, a metric may be a composite metric that combines one or more abovementioned metrics through various suitable ways, such as weighted average or other statistics. In some embodiments, a composite metric may also take the form of a feature vector that

may include multiple dimensions, each dimension corresponding to an individual metric.

(76) The rated historical query records may be used as training samples for the training of one or more machine learning models in the item query engine **330** in model training **650**. For example, the metric values of the rated historical query records may be used as labels of the training samples in training a supervised learning model. In one or more embodiments, a quality metric value may be assigned to each of the historical query records in the representative set that is sampled. The representative set is used to train the item query engine **330** (e.g., one or more machine learning models), such as using a supervised learning technique. For example, the online concierge system **102** uses the item query engine **330** to generate, in a forward propagation of the training, query results using the search phrases of the historical query records in the representative set. The online concierge system **102** compares the query results to the historical query records and the metric values. The online concierge system **102** in turn adjusts, in a back propagation of the training, parameters of the item query engine **330** based on the comparison. For example, the item query engine **330** may be adjusted (automatically, e.g., by the online concierge system **102**) such that it generates more query results with higher quality metric values and fewer query results that have lower quality metric values. Machine learning techniques such as stochastic gradient descent may be used in adjusting the parameters in the back propagation. The forward propagation and back propagation may be repeated iteratively until the training is completed, such as when the machine learning model converges or after a predetermined number of iterations.

#### Example Query Log

(77) FIG. 7 is a conceptual diagram illustrating an example query log that includes a number of historical query records, according to one or more embodiments. The data structure and the fields in the query log are for illustration only. In various embodiments, a query log may be represented in a different format and with different data fields. Also, the entry in the query log may include additional data fields that are not displayed in FIG. 7. A query log may be specific to a warehouse **210** or a physical store. Another warehouse **210** or another physical store may have another query log.

(78) In FIG. 7, the example query log includes multiple entries. Each entry is a historical query record that may be uniquely identified by a search identifier (Search\_ID). The entry may take the form of key-value pairs where the key is the search identifier and the value includes a list of data fields that record the nature of the historical query record. The data fields may include the search phrase, the time stamp of when the query occurred, and the item results. Each historical query record is associated with a search phrase that includes one or more keywords. However, the search phrases are not unique among the historical query records. For example, in the illustration of FIG. 7, the historical query record with the Search\_ID\_1 and the historical query record with the Search\_ID\_N have the same search phrase. The online concierge system **102** may calculate the search frequencies of different search phrases used in various historical query records in a period of time. Some search phrases are more frequently searched by different users than other phrases.

(79) The item results are records of actual items that are selected and ranked by the item query engine **330**. The item results may be a ranked list of items and the items' corresponding unique identifiers. Since the item query engine **330** conducts searches based on various factors and some of those factors may be user-specific or time-sensitive, the item results are often different even for the same search phrase. For example, in the illustration shown in FIG. 7, while the historical query record with the Search\_ID\_1 and the historical query record with the Search\_ID\_N have the same search phrase, the ranked item results for the two query results are different.

#### Example Query Record Sampling

(80) FIG. 8 is a flowchart depicting an example process **800** for sampling historical query records, according to one or more embodiments. In various embodiments, the process **800** includes different or additional steps than those described in conjunction with FIG. 8. Further, in some embodiments, the steps of the process **800** may be performed in different orders than the order described in

conjunction with FIG. 8. The process **800** may be carried out by the online concierge system **102** in various embodiments. In other embodiments, the steps of the process **800** are performed by any suitable online system. While the process **800** is described to be used in association with an item query engine, the process **800** may also be used for other query engines that are not used for item selection, such as a document retrieval engine, a search engine, and a literature search engine, etc. (81) In some embodiments, the online concierge system **102** retrieves **810** a set of historical query records associated with an item query engine **330**. For example, the online concierge system **102** may retrieve the historical query records in a query log **340**. The set of historical query records includes a plurality of search phrases. A historical query record is associated with a search phrase and a list of items returned by the item query engine **330**, as illustrated in FIG. 7. In some embodiments, the set of historical query records may be retrieved from query records that are within a predetermined number of days old. For example, the process **800** may be performed periodically (e.g., every week, every N days) to collect new sets of representative historical query records that may be used as training samples to further train or reinforce one or more machine learning models used in the item query engine **330** periodically.

(82) The online concierge system **102** determines **820** search frequencies for the search phrases included in the set of the historical query records. The online concierge system **102** may sort the historical query records by search frequencies of the search phrases associated with the historical query records. For example, there may be historical query records that share the same search phrases, particularly those for common search phrases such as “milk,” “steak,” etc. The online concierge system **102** may sort each historical query record based on how common the search phrase associated with the query record is used in other historical query records in a given period of time. For example, the search frequency for a search phrase may correspond to the number of historical query records that have the same search phrase.

(83) In some embodiments, the historical query records may first be classified into two or more groups. For some rare search phrases, each of them may have only one or a handful of historical query records that have the same search phrase. The historical query records with those rare search phrases may be referred to as “singleton” query records. The online concierge system **102** may define a threshold search frequency (e.g., smaller or equal to 1, 2, 3, etc.). The historical query records with search phrases whose search frequencies are under the threshold may be classified as singleton queries while other historical query records with more common search phrases may be classified as “non-singleton” or normal queries. While two groups (singleton or non-singleton) are given as examples here, the historical query records may be classified into multiple groups. For example, multiple search frequency threshold levels may be defined. Regardless of whether a historical query record is singleton or not or is classified or not, the historical query records in the set of a given period may be sorted by the search frequencies of the search phrases.

(84) The online concierge system **102** stratifies **830** the set of historical query records into a plurality of bins according to the search frequencies of the search phrases. A bin corresponds to a subset of historical query records whose search phrases' search frequencies are within a range. A bin size may refer to the total number of bins used. A bin size for a particular group may be the total number of bins used in the group. In some embodiments, before stratifying the historical query records into bins, the historical query records are classified as groups based on whether the historical query records are singleton or not (or further classified into multiple groups). The singleton group and the non-singleton group may have the same or different bin sizes (e.g., different numbers of bins within a group).

(85) The overall bin size or bin sizes for different groups may be predetermined or determined dynamically. In one or more embodiments, the online concierge system **102** may determine dynamically a bin size based on a targeted output number (targeted sample size). For example, the sampling process **800** may target a particular number of records (e.g., 2000 records) that should be outputted for evaluation. The online concierge system **102** may sample the historical query records

using a bin size that generates a targeted sample size. In one or more embodiments, the online concierge system **102** may determine the bin size based on the targeted output using a binary search technique.

(86) By way of example, the online concierge system **102** may calculate appropriate bin size  $b$  to yield a targeted sample size  $M$  equal to the desired number of record outputs from a query log of volume  $V$  (e.g., the total number of historical query records in a given group, such as the non-singleton group). The online concierge system **102** may use binary search to find the value of  $b$  that yields the closest sample to  $M$ . In one or more embodiments, the online concierge system **102** may divide  $b$  values into two halves on either side of  $\text{value} = [(\# \text{distinct queries} - 1) - 1] / 2$ . If  $M$  is too low, the online concierge system **102** may select  $b$  in the middle of the lower half. If  $M$  is too high, the online concierge system **102** may select  $b$  in the middle of the upper half. The binary search may be repeated iteratively until the online concierge system **102** finds the largest  $b$  that is a power of 2 and satisfies  $b = V/M$ . The number of records in each bin may be roughly equal to the numbers in other bins.

(87) In some embodiments, the bin sizes may vary depending on the types of search records. For example, the online concierge system **102** determines a threshold for singleton queries. As discussed, a singleton query may refer to a query with a search phrase whose search frequency is lower than (or equal to) the threshold. The online concierge system **102** determines a first bin size for non-singleton queries, such as based on the binary search discussed above. The online concierge system **102** determines a second bin size for the singleton queries. The second bin size can be the same or different from the first bin size. The determination of the bin size for the singleton queries may be based on another algorithm. For example, the online concierge system **102** may recompute the bin size  $b_{\text{sub.s}}$  for singleton queries. The online concierge system **102** may define another tunable parameter called a singleton ratio. In one or more embodiments, the singleton ratio,  $\beta$ , may be defined as the ratio of the singleton volume to the overall query volume ( $V$ ). This parameter may be changed to reduce the bias of sampling towards the singleton group or towards the non-singleton group. The bin size  $b_{\text{sub.s}}$  for singleton queries may be calculated as  $\beta = (1 - \beta)(Q - Q') / (\beta * V')$ . In this equation,  $Q$  is the number of distinct queries in the overall query log.  $Q'$  is the number of distinct queries in the query log so far.  $V$  is the total query log volume and  $V'$  is the total query log volume so far.

(88) The online concierge system **102** stratifies the set of historical query records according to the first bin size and the second bin size. In some embodiments, the first bin size is larger than the second bin size. In some embodiments, the ratio of the first bin size to the second bin size corresponds to the ratio of a first number of the non-singleton queries in the set of historical query records to a second number of the singleton queries in the set of historical query records.

(89) FIG. 9 is a graphical illustration of stratifying the historical query records into different bins, according to one or more embodiments. The actual number of bins and the number of records in each bin are examples only and are simplified from an actual example. In the example illustrated in FIG. 9, the singleton records may have a smaller bin size (e.g., a smaller number of bins) but each bin may hold more historical query records. The non-singleton records may have a higher bin size but each bin may hold fewer historical query records. Each bin may have an equal or a similar number of records. The historical query records are sorted by the search frequencies of the search phrases associated with the records, with the leftmost records having the lowest search frequencies and the rightmost records having the highest search frequencies.

(90) The online concierge system **102** samples **840** the historical query records from the plurality of bins to collect a representative set of historical query records. For example, the representative set of historical query records may be sampled evenly from each bin (1 sample per bin,  $N$  sample per bin). In some embodiments, the sampling may also be done unevenly for some of the bins, depending on the selection criteria. In some embodiments, to get the next incremental sample, the online concierge system **102** may double the sample size  $M$  and recompute the sample using the

above steps in the process **800**. By using the same random seed, half of the sample in the existing pool is the same as the previous sample. This creates an incremental sampling. The representative set is the sampled result and may be used for further evaluation.

(91) In some embodiments, the online concierge system **102** may also select 850 items for each historical query record in the representative set. In some embodiments, the online concierge system **102** may not further select items in the item results. For example, referring briefly back to FIG. 7, the sampled historical query record may include all of the items in the ranked item results. In other embodiments, the online concierge system **102** may further select 850 items for each historical query record in the representative set. Selecting items for a historical query record may be based on different rules, depending on the nature of the items and the number of items in the query result. For example, the online concierge system **102** may use a first rule for selecting sponsored items and a second rule for selecting non-sponsored items. The first and second rules may be different. For example, the online concierge system **102** may select only the top items for the sponsored items whereas the online concierge system **102** may select additional items from top N items for the non-sponsored items. If the total number of items recorded in the returned query result is lower than a number, all of the items may be selected.

(92) In some embodiments, the online concierge system **102** outputs **860** the representative set of historical query records for the rating of the historical query records. For example, the online concierge system **102** may output **860** the representative set to another engine of the online concierge system **102** to automatically rate the historical query records. In some embodiments, the online concierge system **102** may output **860** the representative set to a third party rater through an API. Examples of how the historical query records may be rated and may be used for model training are discussed in FIG. 6 in association with elements **640** and **650**.

(93) Various embodiments allow online systems, such as large-scale database systems, to improve the accuracy of the performance evaluation of their search query engines and training of the search query engines. For example, during the normal course of operation of a search query engine, certain common search phrases are often present more frequently than other search phrases. Without a proper strategy of sampling, the sampled results are often focused on the common search phrases. As a result, the performance evaluation of a search query engine and the corresponding training may bias towards the common search phrases. The bias may adversely affect the performance of the query engine in generating results for more uncommon search phrases as those uncommon search phrases may be underrepresented in the sampled results. The embodiments described address this challenge and provide more balanced sampling approaches in evaluating and training query engines.

(94) Various unbiased sampling processes in accordance with one or more embodiments also improve the speed of training of query engines. A biased sampling approach that focuses on common search phrases generates a training set that is skewed towards certain search phrases. This could result in difficulty in reaching convergence of the machine learning model because the training of the machine learning model using a biased training set could result in the model not being able to reach the convergent state. The improvement in the speed of training using an unbiased sampling process in accordance with one or more embodiments results in more efficient usage of computing resources, reduced processor usage and processing time, reduced consumption of network bandwidth and memory usage, and reduction of power consumption.

(95) Additional Considerations

(96) The foregoing description of the embodiments of the invention has been presented for the purpose of illustration; it is not intended to be exhaustive or to limit the invention to the precise forms disclosed. Persons skilled in the relevant art can appreciate that many modifications and variations are possible in light of the above disclosure.

(97) Some portions of this description describe the embodiments of the invention in terms of algorithms and symbolic representations of operations on information. These algorithmic

descriptions and representations are commonly used by those skilled in the data processing arts to convey the substance of their work effectively to others skilled in the art. These operations, while described functionally, computationally, or logically, are understood to be implemented by computer programs or equivalent electrical circuits, microcode, or the like. Furthermore, it has also proven convenient at times, to refer to these arrangements of operations as modules, without loss of generality. The described operations and their associated modules may be embodied in software, firmware, hardware, or any combinations thereof.

(98) Any of the steps, operations, or processes described herein may be performed or implemented with one or more hardware or software modules, alone or in combination with other devices. In one or more embodiments, a software module is implemented with a computer program product comprising a computer-readable medium containing computer program code, which can be executed by a computer processor for performing any or all of the steps, operations, or processes described.

(99) Embodiments of the invention may also relate to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, and/or it may comprise a computing device selectively activated or reconfigured by a computer program stored in the computer. Such a computer program may be stored in a tangible computer readable storage medium, which include any type of tangible media suitable for storing electronic instructions and coupled to a computer system bus. Furthermore, any computing systems referred to in the specification may include a single processor or may be architectures employing multiple processor designs for increased computing capability.

(100) Embodiments of the invention may also relate to a computer data signal embodied in a carrier wave, where the computer data signal includes any embodiment of a computer program product or other data combination described herein. The computer data signal is a product that is presented in a tangible medium or carrier wave and modulated or otherwise encoded in the carrier wave, which is tangible, and transmitted according to any suitable transmission method.

(101) Finally, the language used in the specification has been principally selected for readability and instructional purposes, and it may not have been selected to delineate or circumscribe the inventive subject matter. It is therefore intended that the scope of the invention be limited not by this detailed description, but rather by any claims that issue on an application based hereon.

Accordingly, the disclosure of the embodiments of the invention is intended to be illustrative, but not limiting, of the scope of the invention, which is set forth in the following claims.

## Claims

1. A method, comprising: at a computer system comprising at least one processor and memory: receiving a plurality of search queries from a plurality of users directed at an item query engine of an online system, each search query including a search phrase used by a user to conduct the search query; generating a training set comprising a plurality of training samples, the plurality of training samples comprising a plurality of search phrases, each training sample corresponding to a search query of the item query engine and comprising a search phrase and a list of items returned by the item query engine; generating a representative set of training samples, wherein generating the representative set of training samples comprises: determining search frequencies of the search phrases used in the training samples, stratifying the set of training samples into a plurality of bins according to the search frequencies of the search phrases, wherein each bin of the plurality of bins includes a subset of training samples, wherein each bin defines a range of numbers of times a search phrase is used by the plurality of users, and sampling the training samples from the plurality of bins to collect the representative set of training samples; and training the item query engine using the representative set of training samples.

2. The method of claim 1, wherein training the item query engine using the representative set of

training samples comprises: initiating a plurality of parameters of the item query engine and a loss function; applying the item query engine to the training samples to generate predictions of labels of the training samples; comparing the predictions of labels of the training samples to the training labels of the training samples based on the loss function; and backpropagating the loss function to adjust the plurality of parameters of the item query engine.

3. The method of claim 1, wherein generating the training set comprises: retrieving, for a training sample, the list of items returned by the item query engine; generating, for each item in list in the training sample, an item embedding based on product data corresponding to the item; and storing item embeddings corresponding to the list of items as part of the training sample.

4. The method of claim 1, wherein generating the training set comprises: retrieving, for a training sample, the user who conducted the search query; generating, for the user in the training sample, a user embedding based on user data that comprises a user profile and a historical purchase; and storing the user embedding as part of the training sample.

5. The method of claim 1, wherein generating the training set comprises: generating, for a training sample, a user embedding corresponding the user who conducted the search query; generating, for an item in the training sample, an item embedding based on product data; and generating, for the training sample, an anchor embedding based on the search query and the user embedding; and generating an item score for the item based on comparing the anchor embedding and the item embedding.

6. The method of claim 1, wherein generating the training set comprises: monitoring the item query engine for a period of time; storing the plurality of search queries received within the period of time as a set of historical query records; and determining the search frequencies of the search phrases in the set of historical query records, wherein a search frequency for a given search phrase is a number of times the given search phrase used by the plurality of users within the period of time.

7. The method of claim 1, further comprising: assigning a quality metric value to each of the training samples in the representative set, resulting in a plurality of quality metric values; and using the representative set and the plurality of quality metric values to train a machine learning model in the item query engine.

8. The method of claim 1, further comprising: dynamically determining a bin size for the plurality of bins through a binary search, wherein sampling the training samples generates a predetermined number of historical query records.

9. The method of claim 1, wherein stratifying the set of training samples into a plurality of bins comprises: determining a threshold for singleton queries, wherein a singleton query is a query with a search phrase whose search frequency is lower than the threshold; determining a first bin size for non-singleton queries; determining a second bin size for the singleton queries, wherein the second bin size is different from the first bin size; and stratifying the set of training samples according to the first bin size and the second bin size.

10. A non-transitory computer-readable medium configured to store code comprising instructions, wherein the instructions, when executed by one or more processors, cause the one or more processors to: receive a plurality of search queries from a plurality of users directed at an item query engine of an online system, each search query including a search phrase used by a user to conduct the search query; generate a training set comprising a plurality of training samples, the plurality of training samples comprising a plurality of search phrases, each training sample corresponding to a search query of the item query engine and comprising a search phrase and a list of items returned by the item query engine; generate a representative set of training samples, wherein generating the representative set of training samples comprises: determining search frequencies of the search phrases used in the training samples, stratifying the set of training samples into a plurality of bins according to the search frequencies of the search phrases, wherein each bin of the plurality of bins includes a subset of training samples, wherein each bin defines a



range of numbers of times a search phrase is used by the plurality of users, and sampling the training samples from the plurality of bins to collect the representative set of training samples; and training the item query engine using the representative set of training samples.

11. The non-transitory computer-readable medium of claim 10, wherein the instruction to train the item query engine using the representative set of training samples comprises instructions to: initiate a plurality of parameters of the item query engine and a loss function; apply the item query engine to the training samples to generate predictions of labels of the training samples; compare the predictions of labels of the training samples to the training labels of the training samples based on the loss function; and backpropagate the loss function to adjust the plurality of parameters of the item query engine.

12. The non-transitory computer-readable medium of claim 10, wherein the instruction to generate the training set comprises instructions to: retrieve, for a training sample, the list of items returned by the item query engine; generate, for each item in list in the training sample, an item embedding based on product data corresponding to the item; and store item embeddings corresponding to the list of items as part of the training sample.

13. The non-transitory computer-readable medium of claim 10, wherein the instruction to generate the training set comprises instructions to: retrieve, for a training sample, the user who conducted the search query; generate, for the user in the training sample, a user embedding based on user data that comprises instructions to a user profile and a historical purchase; and store the user embedding as part of the training sample.

14. The non-transitory computer-readable medium of claim 10, wherein the instruction to generate the training set comprises instructions to: generate, for a training sample, a user embedding corresponding the user who conducted the search query; generate, for an item in the training sample, an item embedding based on product data; and generate, for the training sample, an anchor embedding based on the search query and the user embedding; and generate an item score for the item based on comparing the anchor embedding and the item embedding.

15. The non-transitory computer-readable medium of claim 10, wherein the instruction to generate the training set comprises instructions to: monitor the item query engine for a period of time; store the plurality of search queries received within the period of time as a set of historical query records; and determine the search frequencies of the search phrases in the set of historical query records, wherein a search frequency for a given search phrase is a number of times the given search phrase used by the plurality of users within the period of time.

16. The non-transitory computer-readable medium of claim 10, wherein the instructions, when executed, further cause the one or more processors to: assign a quality metric value to each of the training samples in the representative set, resulting in a plurality of quality metric values; and use the representative set and the plurality of quality metric values to train a machine learning model in the item query engine.

17. The non-transitory computer-readable medium of claim 10, wherein the instructions, when executed, further cause the one or more processors to: dynamically determine a bin size for the plurality of bins through a binary search, wherein sampling the training samples generates a predetermined number of historical query records.

18. The non-transitory computer-readable medium of claim 10, wherein the instruction to stratify the set of training samples into a plurality of bins comprises instructions to: determine a threshold for singleton queries, wherein a singleton query is a query with a search phrase whose search frequency is lower than the threshold; determine a first bin size for non-singleton queries; determine a second bin size for the singleton queries, wherein the second bin size is different from the first bin size; and stratify the set of training samples according to the first bin size and the second bin size.

19. A system comprising: one or more processors; and memory configured to store code comprising instructions, wherein the instructions, when executed by the one or more processors,

cause the one or more processors to: receive a plurality of search queries from a plurality of users directed at an item query engine of an online system, each search query including a search phrase used by a user to conduct the search query; generate a training set comprising a plurality of training samples, the plurality of training samples comprising a plurality of search phrases, each training sample corresponding to a search query of the item query engine and comprising a search phrase and a list of items returned by the item query engine; generate a representative set of training samples, wherein generating the representative set of training samples comprises: determining search frequencies of the search phrases used in the training samples, stratifying the set of training samples into a plurality of bins according to the search frequencies of the search phrases, wherein each bin of the plurality of bins includes a subset of training samples, wherein each bin defines a range of numbers of times a search phrase is used by the plurality of users, and sampling the training samples from the plurality of bins to collect the representative set of training samples; and training the item query engine using the representative set of training samples.

20. The system of claim 19, wherein the instruction to train the item query engine using the representative set of training samples comprises instructions to: initiate a plurality of parameters of the item query engine and a loss function; apply the item query engine to the training samples to generate predictions of labels of the training samples; compare the predictions of labels of the training samples to the training labels of the training samples based on the loss function; and backpropagate the loss function to adjust the plurality of parameters of the item query engine.

---