

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent

12393601

Kind Code

B2

Date of Patent

August 19, 2025

Inventor(s)

Peignier; Timothée et al.

---

### Distributed data processing using embedded hermetic and deterministic language

---

#### Abstract

In one embodiment, a computer system comprises one or more central processing units (CPUs) that are communicatively coupled to a system clock, one or more network interfaces, and one or more database interfaces; digital electronic main memory that is communicatively coupled to the one or more CPUs and storing one or more sequences of stored program instructions which, when executed using the one or more CPUs, cause the one or more CPUs to execute a plurality of different consumer services of a SaaS-based data analytics platform, each of the consumer services hosting an instance of a sandboxed runtime for a hermetic and deterministic programming language; user function storage that is communicatively coupled to one of the database interfaces and storing a plurality of different user functions, each of the user functions having been programmed using the programming language, each of the user functions being stored in association with a reference to a destination table of a destination database; each of the consumer services being programmed to initiate a data ingestion process; load a copy of a user function from the user function storage to the sandboxed runtime that is associated with the particular consumer service; using the sandboxed runtime local to the particular consumer service, execute the user function over records directed to the destination table identified in the reference; filter the records or write new records resulting from the function to the destination table.

---

**Inventors:** Peignier; Timothée (Vancouver, CA), Steel; Edward (Vancouver, CA)

**Applicant:** Treasure Data, Inc. (Mountain View, CA)

**Family ID:** 1000008762196

**Assignee:** Treasure Data, Inc. (Mountain View, CA)

**Appl. No.:** 18/638146

**Filed:** April 17, 2024

#### Prior Publication Data

**Document Identifier**

**Publication Date**

## Related U.S. Application Data

continuation parent-doc US 17815234 20220727 US 12013867 child-doc US 18638146

---

## Publication Classification

**Int. Cl.:** G06F16/25 (20190101); G06F8/30 (20180101); G06F8/70 (20180101)

**U.S. Cl.:**

**CPC** G06F16/252 (20190101); G06F8/31 (20130101); G06F8/70 (20130101);

## Field of Classification Search

**CPC:** G06F (16/252)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11403298	12/2021	Brossard	N/A	N/A
11726976	12/2022	Brossard	707/702	G06F 16/248
2020/0019558	12/2019	Okorafor	N/A	G06F 21/6254

### OTHER PUBLICATIONS

Google/starlark-go, Question: Prevent long-running scripts #160. Nov. 6, 2020 [retrieved Feb. 19, 2025]. Retrieved from the Internet:

<https://web.archive.org/web/20201106232700/https://github.com/google/starlark-go/issues/160> (Year: 2020). cited by examiner

Starlark Language Specification, May 11, 2020 [retrieved on Feb. 15, 2024]. Retrieved from the Internet:

<https://web.archive.org/web/20200511205044/https://github.com/bazelbuild/starlark/blob/master/spec.md>. (Year: 2020), 67 pages. cited by applicant

---

*Primary Examiner:* Andersen; Kristopher

*Attorney, Agent or Firm:* Baker Botts L.L.P.

---

## Background/Summary

**BENEFIT CLAIM (1)** This application claims the benefit under 35 U.S.C. § 120 as a continuation of application Ser. No. 17/815,234, filed Jul. 27, 2022, the entire contents of which are hereby incorporated by reference as if fully set forth herein. The Applicant hereby rescinds any disclaimer of claim scope in the application(s) of which the benefit is claimed and advises the USPTO that the present claims may be broader than any application(s) of which the benefit is claimed.

## COPYRIGHT NOTICE

(1) A portion of the disclosure of this patent document contains material which is subject to copyright protection. The copyright owner has no objection to the facsimile reproduction by anyone of the patent document or the patent disclosure, as it appears in the Patent and Trademark Office patent file or records, but otherwise reserves all copyright or rights whatsoever. ©2021-2022 Treasure Data, Inc.

## TECHNICAL FIELD

(2) One technical field of the present disclosure is relational database systems, including automated transformations of database records.

## BACKGROUND

(3) The approaches described in this section are approaches that could be pursued, but not necessarily approaches that have been previously conceived or pursued. Therefore, unless otherwise indicated, it should not be assumed that any of the approaches described in this section qualify as prior art merely by virtue of their inclusion in this section.

(4) Digital relational databases now are widely used to store, manage, and control the flow of information in business enterprises, with integration into a variety of practical applications. For example, a customer data platform (CDP) is a special-purpose combination of database storage technology and server application software that facilitate storing millions of data signals representing interactions of consumers or end users with merchant websites, social media, digital advertising, and other resources, conducting analysis of the signals, forming audience segments, and other tasks. Supplier of CDPs typically grant access, to enterprise customers, to an instance of a CDP that operates with multi-tenant, siloed data that is subject to security controls.

(5) Customers of CDPs create and store business logic that defines transformations of datasets and other operations on data stored in the CDP. Customers would like to have their business logic applied to their data across all applications of the CDP in a uniform way. However, when business logic that was created for one application is applied to a second, different application, unexpected results can occur. Further, the business logic created for one application can be treated as non-trusted from the perspective of a second application.

(6) One approach to this issue is centralized execution of the data processing logic. However, this approach creates a single point of failure and can lead to resource contention if multiple different applications need to execute the same logic at the same time.

## SUMMARY

(7) The appended claims may serve as a summary of the invention.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) In the drawings:

(2) FIG. 1 illustrates a distributed computer system showing the context of use and principal functional elements with which one embodiment could be implemented.

(3) FIG. 2 illustrates an example process or algorithm that can be programmed to implement one embodiment of executing a user function locally in a consumer service of a customer data platform.

(4) FIG. 3 illustrates an example data flow for data elements and functional elements of an embodiment in a customer data platform (CDP).

(5) FIG. 4 illustrates a process flow or algorithm that can be programmed to centrally manage user functions and execute the user functions at different consumer processes of the CDP using hermetic and deterministic approaches.

(6) FIG. 5 illustrates a computer system with which one embodiment could be implemented.

### DETAILED DESCRIPTION

(7) In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of the present invention. It will be apparent, however, that the present invention may be practiced without these specific details. In other instances, well-known structures and devices are shown in block diagram form in order to avoid unnecessarily obscuring the present invention.

(8) The text of this disclosure, in combination with the drawing figures, is intended to state in prose the algorithms that are necessary to program a computer to implement the claimed inventions, at the same level of detail that is used by people of skill in the arts to which this disclosure pertains to communicate with one another concerning functions to be programmed, inputs, transformations, outputs and other aspects of programming. That is, the level of detail set forth in this disclosure is the same level of detail that persons of skill in the art normally use to communicate with one another to express algorithms to be programmed or the structure and function of programs to implement the inventions claimed herein.

#### 1. General Overview

(9) Centrally managed business logic is distributed to a plurality of different applications of a federated system using a hermetic and deterministic language. The hermetic and deterministic language provides a sandbox to safely execute customer business logic by preventing unsafe access and/or excessive use of resources. The language offers extensive options for data computation and allows virtual routing of logic to different applications. Customers can execute their business logic without constraints but cannot inadvertently or maliciously escape the sandbox. Furthermore, individual applications can change parts of the business logic locally. The creator of the business logic has a single view which elides the many components that are involved.

(10) Embodiments can be applied in the context of data ingestion processes in which increased data processing efficiency is achieved by filtering or transforming records of input datasets before storing records in a destination table of a destination database. Example functions that can be performed during data ingestion include masking, anonymization, removal of sensitive financial or personal information, replacing values representing personal income with a band or range, or excluding data from a sensitive country. Embodiments also facilitate repeated execution of the same logic uniformly using more than one application, process, or consumer service. Safe execution without adverse impact on the platform also is important since user functions could be customer-defined outside the control of an administrator of the platform.

(11) Goals of hermetic and deterministic execution can be achieved by implementing user functions in a programming language having a runtime with those characteristics. In one embodiment, a JAVA implementation of the STARLARK environment is used, but others are possible. In this environment, a user function can define substantive data processing operations but cannot access disk storage or execute network calls; furthermore, the runtime cannot be modified to add these functions. In some embodiments, the runtime is extended to provide REGEX support, and is configured with a maximum number of run iterations to enable shutting down the runtime when an ingestion operation involves so many records that system performance could suffer.

(12) In various embodiments, the disclosure encompasses the subject matter of the following numbered clauses:

(13) 1. A computer system comprising: one or more central processing units (CPUs) that are communicatively coupled to a system clock, one or more network interfaces, and one or more database interfaces; digital electronic main memory that is communicatively coupled to the one or more CPUs and storing one or more sequences of stored program instructions which, when executed using the one or more CPUs, cause the one or more CPUs to execute a plurality of different consumer services of a SaaS-based data analytics platform, each of the consumer services hosting an instance of a sandboxed runtime for a hermetic and deterministic programming language; user function storage that is communicatively coupled to one of the database interfaces and storing a plurality of different user functions, each of the user functions having been

programmed using the programming language, each of the user functions being stored in association with a reference to a destination table of a destination database; each of the consumer services being programmed to initiate a data ingestion process; load a copy of a user function from the user function storage to the sandboxed runtime that is associated with the particular consumer service; using the sandboxed runtime local to the particular consumer service, execute the user function over records directed to the destination table identified in the reference; filter the records or write new records resulting from the function to the destination table.

(14) 2. The computer system of clause 1, the SaaS-based data analytics platform comprising a customer data platform (CDP).

(15) 3. The computer system of clause 1, the programming language comprising STARLARK.

(16) 4. The computer system of clause 1, each of the consumer services being programmed to delete the user function from the particular consumer service or its sandboxed runtime after completing the data ingestion process.

(17) 5. The computer system of clause 1, each of the consumer services being programmed to: receive an identifier of a source dataset for ingestion; read successive records of the source dataset until all records have been read; for each record of the successive records that is read, determine whether a user function is present in the user function storage and has a reference to the destination table, and only when the user function is present and has the reference to the destination table, execute the user function over the each record.

(18) 6. The computer system of clause 5, each of the consumer services being programmed to write a new record resulting from executing the user function only when the new record is valid.

(19) 7. The computer system of clause 1, each of the consumer services being programmed to: asynchronously, with respect to the execution of the particular consumer service, using a plurality of other consumer services among the plurality of different consumer services executing in a SaaS-based data analytics platform, initiate a plurality of other data ingestion processes for a plurality of datasets; as part of each data ingestion process among the plurality of other data ingestion processes, load a second copy of the user function using a second sandboxed runtime that is associated with each other consumer service; using a second sandboxed runtime that is local to each other consumer service, execute, over records of the datasets, the second copy of the user function directed to the destination table identified in the reference and filter the records or write new records resulting from the function to the destination table.

(20) 8. A computer-implemented method comprising: receiving a digitally defined user function that has been coded using a hermetic and deterministic human-readable programming language; storing the user function in a centralized database of user functions, in a table of user functions with a reference to a destination table; at a particular consumer service among a plurality of different consumer services executing in a SaaS-based data analytics platform, initiating a data ingestion process; loading a copy of the user function using a sandboxed runtime that is associated with the particular consumer service; using the sandboxed runtime local to the particular consumer service, executing the user function over records directed to the destination table identified in the reference and filtering the records or writing new records resulting from the function to the destination table.

(21) 9. The method of clause 8, the SaaS-based data analytics platform comprising a customer data platform (CDP).

(22) 10. The method of clause 8, the programming language comprising STARLARK.

(23) 11. The method of clause 8, further comprising deleting the user function from the particular consumer service or its sandboxed runtime after completing the data ingestion process.

(24) 12. The method of clause 8, the particular consumer service being programmed to: receive an identifier of a source dataset for ingestion; read successive records of the source dataset until all records have been read; for each record of the successive records that is read, determine whether a user function is present in the centralized database of user functions and has a reference to the destination table, and only when the user function is present and has the reference to the destination

table, execute the user function over the each record.

(25) 13. The method of clause 11, each of the consumer services being programmed to write a new record resulting from executing the user function only when the new record is valid.

(26) 14. The method of clause 8, further comprising, after completing the data ingestion process, deleting the user function from the particular consumer service or its sandboxed runtime.

(27) 15. The method of clause 8, further comprising: asynchronously, with respect to the execution of the particular consumer service, using a plurality of other consumer services among the plurality of different consumer services executing in a SaaS-based data analytics platform, initiating a plurality of other data ingestion processes for a plurality of datasets; as part of each data ingestion process among the plurality of other data ingestion processes, loading a second copy of the user function using a second sandboxed runtime that is associated with each other consumer service; using a second sandboxed runtime that is local to each other consumer service, executing, over records of the datasets, the second copy of the user function directed to the destination table identified in the reference and filtering the records or writing new records resulting from the function to the destination table.

## 2. Structural & Functional Overview

(28) FIG. 1 illustrates a distributed computer system showing the context of use and principal functional elements with which one embodiment could be implemented.

(29) In an embodiment, a computer system **100** comprises components that are implemented at least partially by hardware at one or more computing devices, such as one or more hardware processors executing stored program instructions stored in one or more memories for performing the functions that are described herein. In other words, all functions described herein are intended to indicate operations that are performed using programming in a special-purpose computer or general-purpose computer, in various embodiments. FIG. 1 illustrates only one of many possible arrangements of components configured to execute the programming described herein. Other arrangements may include fewer or different components, and the division of work between the components may vary depending on the arrangement.

(30) FIG. 1, and the other drawing figures and all of the description and claims in this disclosure, are intended to present, disclose and claim a technical system and technical methods in which specially programmed computers, using a special-purpose distributed computer system design, execute functions that have not been available before to provide a practical application of computing technology to the problem of how to write once, centrally manage, yet deploy business logic to data across all of a plurality of applications of a SaaS-based application server, such as a CDP, in a uniform way. In this manner, the disclosure presents a technical solution to a technical problem, and any interpretation of the disclosure or claims to cover any judicial exception to patent eligibility, such as an abstract idea, mental process, method of organizing human activity or mathematical algorithm, has no support in this disclosure and is erroneous.

(31) In an embodiment, a user computer **102** is coupled directly or indirectly via a data communication network **103** to a data warehouse **106** and a SaaS-based data analytics platform **104**, which is coupled to user function storage **140** and a destination database **150**. For purposes of illustrating a clear example, one user computer **102** is shown, but in practical embodiments there may be any number, up to thousands or millions, depending on the processing power of platform **104**, with which the user computer has a client-server relationship. User computer **102** can comprise a laptop computer, desktop computer, workstation, tablet computer, or mobile computing device and typically hosts an operating system and one or more application programs or apps, including an internet browser.

(32) The SaaS-based data analytics platform **104** comprises a federated set of computer program applications, sets of instructions, functions, methods, or other software elements that interoperate to provide data analytics functions to user computer **102** via an HTTP server that can generate dynamic web pages for consumption using a browser at the user computer. Broadly, the platform

**104** is programmed to implement data ingestion functions, data storage functions, data analysis, and data display. The platform **104** can implement any of multiple different substantive functions or applications. In one embodiment, platform **104** comprises a customer data platform (CDP) such as the TREASURE DATA platform of Treasure Data, Inc. and Treasure Data K.K. Other embodiments could implement medical data analysis, financial data analysis, business data analysis, education data analysis, or other applications.

(33) Network **103** comprises any combination of one or more local area networks (LANs), wide area networks (WANs), campus networks, or internetworks and can include elements of the public internet. Data warehouse **106** comprises one or more directly attached or networked storage devices, storage clusters, or cloud-based storage systems. Data warehouse **106** can be implemented using systems such as TREASURE DATA CLOUD DATA WAREHOUSE, commercially available from Treasure Data, Inc. and Treasure Data K.K. User function storage **140** and destination database **150** can be implemented using tables in the data warehouse **106** or using independent storage devices and/or database systems.

(34) In an embodiment, user computer **102** can be operated to create and store one or more user functions **105** in the user function storage **140**, as further described in other sections herein. The one or more user functions **105** can be stored in text form in a folder, bucket, or other filesystem construct of a file server. For example, AMAZON S3 can be used for large-capacity file storage. Functions typically are marked for read-only access so that the process of FIG. 2 cannot overwrite functions but only read them. Cache storage can be implemented at the filesystem or at the platform **104** to reduce network roundtrips and to read a function from networked persistent storage only when it is accessed for the first time during a job.

(35) The platform **104** comprises, among other elements, a plurality of different consumer services **108, 110, 112**, each having a sandboxed runtime **114** that is capable of reading, using, and/or hosting a local copy of a user function **116**. Each of the consumer services **108, 110, 112** can comprise any program, set of programs, service, microservice, method, function, object or thread that is capable of receiving a dataset and creating, reading, updating, or deleting records of the dataset in conjunction with stored program functional logic to execute a substantive data processing function. Examples of consumer services include data transformation services, data analytics services, and data presentation services.

(36) The platform **104** comprises, in one embodiment, an operating system **162** that implements services and functions to control access to main memory, I/O devices, a network stack to interface to network **103**, and other services. The platform **104** can comprise a system clock **164** that provides a master point of reference for timing of presentation of executable instructions to a CPU of the CDP, memory accesses, I/O accesses, and other functions. Other hardware elements are not shown for clarity but can include those shown in FIG. 5. The platform **104** may be coupled to or integrated with a file system **142** that provides flat file storage of system programs, copies of pages swapped out from main memory, configuration files, bootstrap loader files, and other system support files. Platform **104** can be implemented using one or more physical computers or virtual machines, with network attached storage or virtual storage, in any of a computing cluster, private data center, public cloud data center, or other virtual computing facility. Examples of public cloud data centers that could be used for virtual machine implementations of platform **104** include AMAZON AWS and related services, MICROSOFT AZURE, GOOGLE CLOUD, etc.

(37) In an embodiment, input from a user computer **102** can specify logical operations, expressed in pseudocode, to process events in a stream of result data from other operations or in a batch, and organized in a form of a script file denoted as the user function **105**. The logical operations can be expressed in program code and/or a database query language, such as Structured Query Language (SQL). In an embodiment, the platform provides a single entry point that can receive either batch-based or stream-based transformations using a sandboxed language. Functions are written once but can be executed repeatedly by different consumers or applications in a federated, multi-tenant CDP

or other applied database system.

(38) In an embodiment, the user function **105** is coded using the STARLARK language.

STARLARK is an open-source system for which program code is available, at the time of this writing, in the path, folder, or directory/bazelbuild/starlark of the GITHUB service at the domain github.com and accessible via the HTTP protocol from a browser. A JAVA runtime is available, at the time of this writing, in the path, folder, or directory/bazelbuild/bazel/trec/master/src/maj/java/net/starlark/java of the GITHUB service at the domain github.com and accessible via the HTTP protocol from a browser. In an embodiment, each sandboxed runtime **114** of platform **104** is a copy of the STARLARK runtime module and is capable of executing a copy of the user function **105**, which is locally stored and used in the form of users functions **116**. Thus, while user function **105** can be created once and centrally stored via user function storage **140**, any of the consumer services **108**, **110**, **112** can, on demand, load a copy of user function **105** from user function storage **140** into memory associated with that consumer service, thus creating and storing user function **116**, and locally execute the user function under the control of sandboxed runtime **114**. When a user function **116** completes execution, it returns any result and/or an exception or error code to the consumer service that hosts it, e.g., consumer service **112**.

(39) Thus, the term “sandboxed” in FIG. **1** can refer, in one embodiment or example, to the attributes of an executable program that is capable of reading or loading, and interpreting or executing, other programs, while enforcing a set of access controls on the code that is interpreted and executed. One set of access controls could be that customer logic in a user function **116** cannot access the file system **142**, data warehouse **106** or destination database **150** or other file storage to which the CDP **140** is connected, cannot access network **103**, and cannot access the system clock **164** of the CDP, so the untrusted code is safe to execute.

(40) As an example, the STARLARK runtime offers the attributes of deterministic execution and hermetic execution. With deterministic execution, executing the same code twice always yields the same results. With hermetic execution, customer logic cannot access the file system, network, or system clock, and therefore untrusted code is safe to execute. For example, customer logic that attempts to implement recursion, appending to an external list, or loading external scripts will fail, as the STARLARK runtime does not permit these operations. In some embodiments, functions are limited to a specified maximum number of steps, to prevent the use of excess resources when jobs are running. For example, in various embodiments, 50 to 500 steps can be used as limits.

(41) FIG. **4** illustrates a process flow or algorithm that can be programmed to centrally manage user functions and execute the user functions at different consumer processes of the CDP using hermetic and deterministic approaches. At block **402**, the process of FIG. **4** is programmed to receive a digitally defined user function that has been coded using a hermetic and deterministic human-readable programming language. The user function can be prepared using a text editor, code editor, source code development environment, or other programming tool, using any of several programming languages, including but not limited to STARLARK, as further described in other sections herein. After the user function has been written and digitally stored in a digital storage device, block **402** can execute to automatically read or retrieve a copy of the user function from a storage location.

(42) At block **404**, the process of FIG. **4** is programmed to store the user function in a centralized database of user functions, in a table of user functions with a reference to a destination table. For example, the user function **105** (FIG. **1**) can be stored in user function storage **140**, which can be formatted according to a table schema that includes a table with rows identifying user functions and one or more column attributes that identify one or more destination tables to which each user function of a row applies.

(43) At block **406**, the process of FIG. **4** is programmed to initiate, at a particular consumer process among a plurality of different consumer processes executing in a customer data platform (CDP), a



data ingestion process. For example, block **406** can comprise a particular consumer process of the CDP, among a plurality of different consumer processes that are then currently executing in the CDP, programmatically initiating a data ingestion process. Block **406** can execute multiple times asynchronously with respect to multiple different consumer processes of the CDP. That is, the CDP can be programmed as a multi-threaded process in which a plurality of different applications, functions, methods, or other executable units are individually executing at once, on different execution paths. Any such application, function, method, or unit can comprise a consumer process that independently initiates a data ingestion process to obtain or receive data useful for a larger task for which the application, function, method, or unit is programmed.

(44) In an embodiment, each such application, function, method, or unit includes, or can load and run, a copy or instance of a sandboxed runtime that constrains the operation of the user function or other functions. At block **408**, the process of FIG. 4 is programmed to load a copy of the user function using the sandboxed runtime that is associated with the particular consumer process.

(45) At block **410**, the process of FIG. 4 is programmed to execute, using the sandboxed runtime local to the particular consumer process, the user function over records directed to the destination table identified in the reference. Such localized, sandboxed execution of the user function can, in various embodiments of the user function, result in filtering records and/or creating new records for the destination table that is identified in user function storage for the user function. Thus, block **412** represents multiple different cases in which the process of FIG. 4 is programmed to filter the records or write new records resulting from the function to the destination table, depending on the substantive content of the function.

(46) At block **414**, the process of FIG. 4 is programmed to delete the user function from the particular consumer process or its sandboxed runtime after completing the data ingestion process. Block **414** can represent reclaiming memory for a copy of the user function that was received and executed local to a particular consumer process. The operation of block **414** does not affect other centralized storage of the user function, such as user function storage **140**. Consequently, at other times after block **414**, any consumer process of the CDP can load the user function again and execute it locally using an instance of the sandboxed runtime that the consumer process manages.

(47) FIG. 2 illustrates an example process or algorithm that can be programmed to implement one embodiment of executing a user function locally in a consumer service of a customer data platform. FIG. 2 and each other flow diagram herein is intended as an illustration at the functional level at which skilled persons, in the art to which this disclosure pertains, communicate with one another to describe and implement algorithms using programming. The flow diagrams are not intended to illustrate every instruction, method object or sub-step that would be needed to program every aspect of a working program, but are provided at the same functional level of illustration that is normally used at the high level of skill in this art to communicate the basis of developing working programs.

(48) At block **202**, the process of FIG. 2 initiates a data ingestion process. For example, block **202** may specify that the process of FIG. 2 generally begins when an end user, system, or script initiates a data ingestion process using a native data ingestion function of a customer data platform. The CDP of TREASURE DATA, for example, implements a Data Ingestion API that a program can call to initiate intake of external data into the databases and table schemas of the CDP. Such a call typically specifies at least a source dataset to be ingested and a destination table.

(49) During data ingestion, at block **204**, the process of FIG. 2 is programmed to read a record of a database table, typically a source table from which data is obtained. At block **206**, the process of FIG. 2 is programmed to test whether sandboxed business logic is enabled for the record that was read at block **204**. Enablement, at block **206**, can be specified using a column attribute of the record, effectively signaling that other steps of FIG. 2 should execute to determine if a user function has been defined.

(50) If the test of block **206** is NO or FALSE or the equivalent, then control transfers to block **208**

at which the process of FIG. 2 is programmed to test whether the record read at block **204** is the last record in the dataset to be ingested. If the test of block **208** is YES or TRUE or the equivalent, then control transfer to block **210** at which the process of FIG. 2 is programmed to end operation or return control to a calling process.

(51) If the test of block **206** is YES or TRUE or the equivalent, then control transfers to block **212** in which the process of FIG. 2 is programmed to test whether a user function is present for a destination table that is the target of the data ingestion process. Block **212** can be programmed to search user function storage **140** (FIG. 1) to identify, in a table of user functions or metadata relating to multiple stored user functions such as the user function **105**, whether a specific user function has been defined for the destination table that was identified in the data ingestion request or API call.

(52) If the test of block **212** is NO or FALSE or the equivalent, then the process of FIG. 2 is programmed to transfer control to block **208**, which has been discussed above. If block **212** is YES or TRUE or the equivalent, then control transfers to block **214** at which the user function is executed over the record of the dataset that had been read at block **204**. The user function may or may not result in returning a new, transformed record. At block **216**, the process of FIG. 2 is programmed to test whether a new record is returned; if not, then the effect of the user function is to filter the current record, so a write operation to a destination table is not required, as shown in block **218**. After block **218**, control can transfer to block **208**.

(53) If the test of block **216** is YES or TRUE or the equivalent, then control transfers to block **220** at which the process of FIG. 2 is programmed to test whether the new record is valid. Validity checking at block **220** can comprise determining that a format of the new record is correct, for example. If the test of block **220** is YES or TRUE or the equivalent, then control transfers to block **222** at which the process of FIG. 2 is programmed to write the new record to persistent storage, for example, in the destination table that was specified in the API call to initiate data ingestion.

(54) Optionally, the process of FIG. 2 is programmed to perform write error checking at block **224**, block **226** and drop the new record if the new record did not write successfully to the persistent storage and/or destination table. During the execution of FIG. 2, records from JAVA objects are mapped to STARLARK objects, and the converse, in main memory.

(55) FIG. 3 illustrates an example data flow for data elements and functional elements of an embodiment in a customer data platform (CDP). In an example data flow **300**, any number of browsers or mobile apps **302**, or other sources, can post or write, to an ingress streaming connector **304** that is adapted to interoperate with the browsers or mobile apps, instructions to initiate a data ingestion operation. A connector **304** can be programmed to write a call to a streaming data ingestion API **306** to initiate an ingestion operation. As data records are ingested via operation of the API **306**, records arrive at and are stored in an ingestion buffer **308**. Function operators can be embedded either using a KVS consumer **310** or ADL consumer **316** and read using a function API **312**, after which a function is executed against records in the buffer **308** and output records are written to a destination database **320**, ADL **318**, or provided to a workflow for downstream processing. The workflow can be programmed, for example, to batch write new records to the ADL using a compatible API **312**.

(56) Thus, in an embodiment, a pipeline or workflow of FIG. 3 can be structured to execute generally as follows. A data ingestion process initiates at a node, which can be any consumer service. Each node supports a table of functions, which can be stored in a file of a flat filesystem subject to access keys or other security controls. For each record processed in ingestion, inspect the table to determine if a function is present; if so, run the function on the data. Each function further specifies a destination table for new output records. An application programming interface supports operations to write a function, upload it, and download it for execution. Embodiments can implement checks or security controls based on account identifiers and keys for the filesystem.

(57) Functions are stored centrally but loaded and cached in main memory of each node. Functions

are compact. “User functions” and central store is “Functions API”. Nodes call the FAPI to obtain functions. Functions stored per target table for the data. Run a data ingestion pipeline. Call the FAPI within the pipeline, The first time that a destination table is used, a node will call the API to obtain and cache the functions.

(58) TABLE 1 illustrates logic configured to process each record on a stream:

(59) TABLE-US-00001 TABLE 1 STREAM-BASED EXTENSION def process\_record(record)  
if record.get(“td:c360:operation”) != None:      adl(record) # Send to Active Data Layer (ADL)  
with changes      return # Do not send to main storage      return record # Send to main storage

(60) In some embodiments, the logic of TABLE 1 could be linked to a workflow stream-based operator as in TABLE 2.

(61) TABLE-US-00002 TABLE 2 STREAM-BASED EXTENSION +filter\_stream\_records:  
stream>: apply.stl # or apply.sql      source: database/table

(62) Logic like that of the examples in the tables can implement filter operations, conversion operations, masking operations, or enrichment operations on records of tables. Using embodiments, customers are not required to learn technical details of batch and stream contexts, facilitating easier reuse of logic in stream and batch applications. For example, assume that a particular customer has selected the following business requirements for a database table titled “web\_store/cart\_events” and representing online digital shopping cart activities of browsers or:

(63) 1. All shopping cart activity is useful for analysis except for events representing a change in a quantity of an item in the cart.

(64) 2. Purchase events should not be persistently stored.

(65) 3. Email addresses should be obfuscated in all storage.

(66) 4. Credit card details should not be persistently stored.

(67) 5. A “warehouse code” supplied by a front-end system should be replaced.

(68) 6. For certain operations, only fields for SKU, date, and type are useful.

(69) These requirements can be implemented in the user function set forth in TABLE 3.

(70) TABLE-US-00003 TABLE 3 SHOPPING CART EVENT RECORD PROCESSING

WAREHOUSE\_CODES = {      “virginia”: “us01”,      “tokyo”: “ap01”,      “eu”: “eu01”,      “korea”:  
“ap02” } def process\_record(record):      record[“warehouse\_code”]  
WAREHOUSE\_CODES.get(record.pop(“warehouse”), “us01”)      record[“email”] =  
sha256(record[“email”])      record.pop(“credit\_card”)      if record[“type”] !=  
‘QUANTITY\_CHANGE’:      kvs(record)      adl({k: v for k, v in record if k in [“sku”, “time”,  
“type”]})      if record[“type”] == ‘PURCHASE’:      return      return record

(71) In this example, returning “None” means that the record will not be persistently stored or provided in the default output. sha256( ) is a function that the STARLARK runtime implements as a built-in function, but a different implementation of SHA-256 digital signature or hash computation also could be used. adl( ) and kvs( ) can be programmed as extension functions in Java that extend STARLARK and are provided to the runtime; they can be programmed to send a duplicate of the input record to side outputs; for example, adl( ) sends output to Treasure Data's Active Data Layer (ADL) or a key-value store (KVS) for mass storage. A host application also can receive these outputs and process them as needed. The functions also apply validation checks, and discard events that are not valid for ADL or KVS operations. A function of this kind can be written and stored in a single storage location and managed by an API having calls to read and supply copies of functions, but execution is distributed; a plurality of different consumer applications load and run the user function for each destination table.

(72) After execution of the example of TABLE 4, the records shown in TABLE 4 would be sent to a data ingestion API from a web server via a JAVASCRIPT SDK:

(73) TABLE-US-00004 TABLE 4 SHOPPING CART EVENT OUTPUT RECORDS {“time”:  
1642788769, “uid”: 1, “type”: “ADD”, “warehouse”: “tokyo”, “email”:  
“happy.customer@home.com”, “sku”: “5550001”} {“time”: 1642788799, “uid”: 1, “type”:

“QUANTITY\_CHANGE”, “warehouse”: “tokyo”, “email”: “happy.customer@home.com”, “sku”: “5550001”, “old\_value”: 1, “new\_value”: 10} {“time”: 1642788969, “uid”: 1, “type”: “PURCHASE”, “email”: “happy.customer@home.com”, “sku”: [“5550001”], “credit\_card”: “4111111111111111”}

(74) After ingestion, the data is written to a stream buffer. Other elements of FIG. 3, such as KVS Replicator, Ingest Aggregator, and Event Dispatcher can read the data in streaming fashion as consumers. For example, an Ingest Aggregator process can be programmed to ingest events by destination table and persistently store the events; it also can run functions on the events before aggregation. In an embodiment, a functions operator would translate the foregoing records to those shown in TABLE 5:

(75) TABLE-US-00005 TABLE 5 AGGREGATED EVENT RECORDS {“time”: 1642788769, “uid”: 1, “type”: “ADD”, “warehouse”: “ap01”, “email”: “ba4bcb21fc92d062d8ae31f8709517d035b0aea973662845a8923d10904c293f”, “sku”: “5550001”} {“time”: 1642788799, “uid”: 1, “type”: “QUANTITY\_CHANGE”, “warehouse”: “ap01”, “email”: “ba4bcb21fc92d062d8ae31f8709517d035b0aea973662845a8923d10904c293f”, “sku”: “5550001”, “old\_value”: 1, “new\_value”: 10}

(76) The Event Dispatcher process can be programmed to route and filter incoming records for ADL as well as other consumers; the process can be implemented as a multi-tenant system that routes events to zero or more downstream single-tenant systems. A functions operator can be implemented to transform and filter records. If no function is configured for a dataset, then the function is not emitted to the ADL output or downstream. Thus, normal usage is to add a function for each ADL user/destination table that calls the adl (record) function. The Event Dispatcher also can use functions as a way to route events to destinations; for example, STARLARK code can be generated from a static routing configuration followed by computing a cluster/function destination field. In the following example, the user is an ADL user, and the function is for data processing; the events of TABLE 6 would be sent downstream:

(77) TABLE-US-00006 TABLE 6 ADL EVENT EXAMPLE {“time”: 1642788769, “type”: “ADD”, “sku”: “5550001”} {“time”: 1642788799, “type”: “QUANTITY\_CHANGE”, “sku”: “5550001”} {“time”: 1642788969, “type”: “PURCHASE”, “sku”: “5550001”}

### 3. Implementation Example—Hardware Overview

(78) According to one embodiment, the techniques described herein are implemented by at least one computing device. The techniques may be implemented in whole or in part using a combination of at least one server computer and/or other computing devices that are coupled using a network, such as a packet data network. The computing devices may be hard-wired to perform the techniques, or may include digital electronic devices such as at least one application-specific integrated circuit (ASIC) or field programmable gate array (FPGA) that is persistently programmed to perform the techniques, or may include at least one general purpose hardware processor programmed to perform the techniques pursuant to program instructions in firmware, memory, other storage, or a combination. Such computing devices may also combine custom hard-wired logic, ASICs, or FPGAs with custom programming to accomplish the described techniques. The computing devices may be server computers, workstations, personal computers, portable computer systems, handheld devices, mobile computing devices, wearable devices, body mounted or implantable devices, smartphones, smart appliances, internetworking devices, autonomous or semi-autonomous devices such as robots or unmanned ground or aerial vehicles, any other electronic device that incorporates hard-wired and/or program logic to implement the described techniques, one or more virtual computing machines or instances in a data center, and/or a network of server computers and/or personal computers.

(79) FIG. 5 is a block diagram that illustrates an example computer system with which an embodiment may be implemented. In the example of FIG. 5, a computer system 500 and instructions for implementing the disclosed technologies in hardware, software, or a combination of

hardware and software, are represented schematically, for example as boxes and circles, at the same level of detail that is commonly used by persons of ordinary skill in the art to which this disclosure pertains for communicating about computer architecture and computer systems implementations. (80) Computer system **500** includes an input/output (I/O) subsystem **502** which may include a bus and/or other communication mechanism(s) for communicating information and/or instructions between the components of the computer system **500** over electronic signal paths. The I/O subsystem **502** may include an I/O controller, a memory controller and at least one I/O port. The electronic signal paths are represented schematically in the drawings, for example as lines, unidirectional arrows, or bidirectional arrows.

(81) At least one hardware processor **504** is coupled to I/O subsystem **502** for processing information and instructions. Hardware processor **504** may include, for example, a general-purpose microprocessor or microcontroller and/or a special-purpose microprocessor such as an embedded system or a graphics processing unit (GPU) or a digital signal processor or ARM processor. Processor **504** may comprise an integrated arithmetic logic unit (ALU) or may be coupled to a separate ALU.

(82) Computer system **500** includes one or more units of memory **506**, such as a main memory, which is coupled to I/O subsystem **502** for electronically digitally storing data and instructions to be executed by processor **504**. Memory **506** may include volatile memory such as various forms of random-access memory (RAM) or other dynamic storage device. Memory **506** also may be used for storing temporary variables or other intermediate information during execution of instructions to be executed by processor **504**. Such instructions, when stored in non-transitory computer-readable storage media accessible to processor **504**, can render computer system **500** into a special-purpose machine that is customized to perform the operations specified in the instructions.

(83) Computer system **500** further includes non-volatile memory such as read only memory (ROM) **508** or other static storage device coupled to I/O subsystem **502** for storing information and instructions for processor **504**. The ROM **508** may include various forms of programmable ROM (PROM) such as erasable PROM (EPROM) or electrically erasable PROM (EEPROM). A unit of persistent storage **510** may include various forms of non-volatile RAM (NVRAM), such as FLASH memory, or solid-state storage, magnetic disk or optical disk such as CD-ROM or DVD-ROM and may be coupled to I/O subsystem **502** for storing information and instructions. Storage **510** is an example of a non-transitory computer-readable medium that may be used to store instructions and data which when executed by the processor **504** cause performing computer-implemented methods to execute the techniques herein.

(84) The instructions in memory **506**, ROM **508** or storage **510** may comprise one or more sets of instructions that are organized as modules, methods, objects, functions, routines, or calls. The instructions may be organized as one or more computer programs, operating system services, or application programs including mobile apps. The instructions may comprise an operating system and/or system software; one or more libraries to support multimedia, programming or other functions; data protocol instructions or stacks to implement TCP/IP, HTTP or other communication protocols; file format processing instructions to parse or render files coded using HTML, XML, JPEG, MPEG or PNG; user interface instructions to render or interpret commands for a graphical user interface (GUI), command-line interface or text user interface; application software such as an office suite, internet access applications, design and manufacturing applications, graphics applications, audio applications, software engineering applications, educational applications, games or miscellaneous applications. The instructions may implement a web server, web application server or web client. The instructions may be organized as a presentation layer, application layer and data storage layer such as a relational database system using structured query language (SQL) or no SQL, an object store, a graph database, a flat file system or other data storage.

(85) Computer system **500** may be coupled via I/O subsystem **502** to at least one output device **512**. In one embodiment, output device **512** is a digital computer display. Examples of a display

that may be used in various embodiments include a touch screen display or a light-emitting diode (LED) display or a liquid crystal display (LCD) or an e-paper display. Computer system **500** may include other type(s) of output devices **512**, alternatively or in addition to a display device. Examples of other output devices **512** include printers, ticket printers, plotters, projectors, sound cards or video cards, speakers, buzzers or piezoelectric devices or other audible devices, lamps or LED or LCD indicators, haptic devices, actuators or servos.

(86) At least one input device **514** is coupled to I/O subsystem **502** for communicating signals, data, command selections or gestures to processor **504**. Examples of input devices **514** include touch screens, microphones, still and video digital cameras, alphanumeric and other keys, keypads, keyboards, graphics tablets, image scanners, joysticks, clocks, switches, buttons, dials, slides, and/or various types of sensors such as force sensors, motion sensors, heat sensors, accelerometers, gyroscopes, and inertial measurement unit (IMU) sensors and/or various types of transceivers such as wireless, such as cellular or Wi-Fi, radio frequency (RF) or infrared (IR) transceivers and Global Positioning System (GPS) transceivers.

(87) Another type of input device is a control device **516**, which may perform cursor control or other automated control functions such as navigation in a graphical interface on a display screen, alternatively or in addition to input functions. Control device **516** may be a touchpad, a mouse, a trackball, or cursor direction keys for communicating direction information and command selections to processor **504** and for controlling cursor movement on display **512**. The input device may have at least two degrees of freedom in two axes, a first axis (e.g., x) and a second axis (e.g., y), that allows the device to specify positions in a plane. Another type of input device is a wired, wireless, or optical control device such as a joystick, wand, console, steering wheel, pedal, gearshift mechanism or other type of control device. An input device **514** may include a combination of multiple different input devices, such as a video camera and a depth sensor.

(88) In another embodiment, computer system **500** may comprise an internet of things (IoT) device in which one or more of the output device **512**, input device **514**, and control device **516** are omitted. Or, in such an embodiment, the input device **514** may comprise one or more cameras, motion detectors, thermometers, microphones, seismic detectors, other sensors or detectors, measurement devices or encoders and the output device **512** may comprise a special-purpose display such as a single-line LED or LCD display, one or more indicators, a display panel, a meter, a valve, a solenoid, an actuator or a servo.

(89) When computer system **500** is a mobile computing device, input device **514** may comprise a global positioning system (GPS) receiver coupled to a GPS module that is capable of triangulating to a plurality of GPS satellites, determining and generating geo-location or position data such as latitude-longitude values for a geophysical location of the computer system **500**. Output device **512** may include hardware, software, firmware and interfaces for generating position reporting packets, notifications, pulse or heartbeat signals, or other recurring data transmissions that specify a position of the computer system **500**, alone or in combination with other application-specific data, directed toward host **524** or server **530**.

(90) Computer system **500** may implement the techniques described herein using customized hard-wired logic, at least one ASIC or FPGA, firmware and/or program instructions or logic which when loaded and used or executed in combination with the computer system causes or programs the computer system to operate as a special-purpose machine. According to one embodiment, the techniques herein are performed by computer system **500** in response to processor **504** executing at least one sequence of at least one instruction contained in main memory **506**. Such instructions may be read into main memory **506** from another storage medium, such as storage **510**. Execution of the sequences of instructions contained in main memory **506** causes processor **504** to perform the process steps described herein. In alternative embodiments, hard-wired circuitry may be used in place of or in combination with software instructions.

(91) The term “storage media” as used herein refers to any non-transitory media that store data

and/or instructions that cause a machine to operation in a specific fashion. Such storage media may comprise non-volatile media and/or volatile media. Non-volatile media includes, for example, optical or magnetic disks, such as storage **510**. Volatile media includes dynamic memory, such as memory **506**. Common forms of storage media include, for example, a hard disk, solid state drive, flash drive, magnetic data storage medium, any optical or physical data storage medium, memory chip, or the like.

(92) Storage media is distinct from but may be used in conjunction with transmission media. Transmission media participates in transferring information between storage media. For example, transmission media includes coaxial cables, copper wire and fiber optics, including the wires that comprise a bus of I/O subsystem **502**. Transmission media can also take the form of acoustic or light waves, such as those generated during radio-wave and infra-red data communications.

(93) Various forms of media may be involved in carrying at least one sequence of at least one instruction to processor **504** for execution. For example, the instructions may initially be carried on a magnetic disk or solid-state drive of a remote computer. The remote computer can load the instructions into its dynamic memory and send the instructions over a communication link such as a fiber optic or coaxial cable or telephone line using a modem. A modem or router local to computer system **500** can receive the data on the communication link and convert the data to a format that can be read by computer system **500**. For instance, a receiver such as a radio frequency antenna or an infrared detector can receive the data carried in a wireless or optical signal and appropriate circuitry can provide the data to I/O subsystem **502** such as place the data on a bus. I/O subsystem **502** carries the data to memory **506**, from which processor **504** retrieves and executes the instructions. The instructions received by memory **506** may optionally be stored on storage **510** either before or after execution by processor **504**.

(94) Computer system **500** also includes a communication interface **518** coupled to bus **502**. Communication interface **518** provides a two-way data communication coupling to network link(s) **520** that are directly or indirectly connected to at least one communication networks, such as a network **522** or a public or private cloud on the Internet. For example, communication interface **518** may be an Ethernet networking interface, integrated-services digital network (ISDN) card, cable modem, satellite modem, or a modem to provide a data communication connection to a corresponding type of communications line, for example an Ethernet cable or a metal cable of any kind or a fiber-optic line or a telephone line. Network **522** broadly represents a local area network (LAN), wide-area network (WAN), campus network, internetwork or any combination thereof. Communication interface **518** may comprise a LAN card to provide a data communication connection to a compatible LAN, or a cellular radiotelephone interface that is wired to send or receive cellular data according to cellular radiotelephone wireless networking standards, or a satellite radio interface that is wired to send or receive digital data according to satellite wireless networking standards. In any such implementation, communication interface **518** sends and receives electrical, electromagnetic or optical signals over signal paths that carry digital data streams representing various types of information.

(95) Network link **520** typically provides electrical, electromagnetic, or optical data communication directly or through at least one network to other data devices, using, for example, satellite, cellular, Wi-Fi, or BLUETOOTH technology. For example, network link **520** may provide a connection through a network **522** to a host computer **524**.

(96) Furthermore, network link **520** may provide a connection through network **522** or to other computing devices via internetworking devices and/or computers that are operated by an Internet Service Provider (ISP) **526**. ISP **526** provides data communication services through a world-wide packet data communication network represented as internet **528**. A server computer **530** may be coupled to internet **528**. Server **530** broadly represents any computer, data center, virtual machine or virtual computing instance with or without a hypervisor, or computer executing a containerized program system such as DOCKER or KUBERNETES. Server **530** may represent an electronic

digital service that is implemented using more than one computer or instance and that is accessed and used by transmitting web services requests, uniform resource locator (URL) strings with parameters in HTTP payloads, API calls, app services calls, or other service calls. Computer system **500** and server **530** may form elements of a distributed computing system that includes other computers, a processing cluster, server farm or other organization of computers that cooperate to perform tasks or execute applications or services. Server **530** may comprise one or more sets of instructions that are organized as modules, methods, objects, functions, routines, or calls. The instructions may be organized as one or more computer programs, operating system services, or application programs including mobile apps. The instructions may comprise an operating system and/or system software; one or more libraries to support multimedia, programming or other functions; data protocol instructions or stacks to implement TCP/IP, HTTP or other communication protocols; file format processing instructions to parse or render files coded using HTML, XML, JPEG, MPEG or PNG; user interface instructions to render or interpret commands for a graphical user interface (GUI), command-line interface or text user interface; application software such as an office suite, internet access applications, design and manufacturing applications, graphics applications, audio applications, software engineering applications, educational applications, games or miscellaneous applications. Server **530** may comprise a web application server that hosts a presentation layer, application layer and data storage layer such as a relational database system using structured query language (SQL) or no SQL, an object store, a graph database, a flat file system or other data storage.

(97) Computer system **500** can send messages and receive data and instructions, including program code, through the network(s), network link **520** and communication interface **518**. In the Internet example, a server **530** might transmit a requested code for an application program through Internet **528**, ISP **526**, local network **522** and communication interface **518**. The received code may be executed by processor **504** as it is received, and/or stored in storage **510**, or other non-volatile storage for later execution.

(98) The execution of instructions as described in this section may implement a process in the form of an instance of a computer program that is being executed, and consisting of program code and its current activity. Depending on the operating system (OS), a process may be made up of multiple threads of execution that execute instructions concurrently. In this context, a computer program is a passive collection of instructions, while a process may be the actual execution of those instructions. Several processes may be associated with the same program; for example, opening up several instances of the same program often means more than one process is being executed. Multitasking may be implemented to allow multiple processes to share processor **504**. While each processor **504** or core of the processor executes a single task at a time, computer system **500** may be programmed to implement multitasking to allow each processor to switch between tasks that are being executed without having to wait for each task to finish. In an embodiment, switches may be performed when tasks perform input/output operations, when a task indicates that it can be switched, or on hardware interrupts. Time-sharing may be implemented to allow fast response for interactive user applications by rapidly performing context switches to provide the appearance of concurrent execution of multiple processes simultaneously. In an embodiment, for security and reliability, an operating system may prevent direct communication between independent processes, providing strictly mediated and controlled inter-process communication functionality.

(99) In the foregoing specification, embodiments of the invention have been described with reference to numerous specific details that may vary from implementation to implementation. The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense. The sole and exclusive indicator of the scope of the invention, and what is intended by the applicants to be the scope of the invention, is the literal and equivalent scope of the set of claims that issue from this application, in the specific form in which such claims issue, including any subsequent correction.



## Claims

1. A computer system comprising: one or more central processing units (CPUs) that are communicatively coupled to a system clock, one or more network interfaces, and one or more database interfaces; digital electronic main memory that is communicatively coupled to the one or more CPUs and storing one or more sequences of stored program instructions which, when executed using the one or more CPUs, cause the one or more CPUs to execute a plurality of different consumer services of a SaaS-based data analytics platform, each of the consumer services hosting an instance of a sandboxed runtime for a hermetic and deterministic programming language; user function storage that is communicatively coupled to one of the database interfaces and storing a plurality of different user functions, each of the user functions having been programmed using the hermetic and deterministic programming language, the hermetic and deterministic programming language preventing each of the user functions from accessing a file system, the one or more network interfaces, and the system clock, each of the user functions being stored in association with a reference to a destination table of a destination database; each of the consumer services being programmed to initiate a data ingestion process, load a copy of a user function from the user function storage to the sandboxed runtime that is associated with a particular consumer service from the plurality of different consumer services, using the sandboxed runtime local to the particular consumer service, execute the user function over records directed to the destination table identified in the reference, and filter the records or write new records resulting from the function to the destination table; the consumer services being further programmed to receive an identifier of a source dataset for ingestion; read successive records of the source dataset until all records have been read; and for each record of the successive records that is read, determine whether a user function is present in the user function storage and has a reference to the destination table, and only when the user function is present and has the reference to the destination table, execute the user function over the each record.
2. The computer system of claim 1, the SaaS-based data analytics platform comprising a customer data platform (CDP).
3. The computer system of claim 1, each of the consumer services being programmed to delete the user function from the particular consumer service or its sandboxed runtime after completing the data ingestion process.
4. The computer system of claim 1, each of the consumer services being programmed to write a new record resulting from executing the user function only when the new record is valid.
5. The computer system of claim 1, each of the consumer services being programmed to: asynchronously, with respect to the execution of the particular consumer service, using a plurality of other consumer services among the plurality of different consumer services executing in the SaaS-based data analytics platform, initiate a plurality of other data ingestion processes for a plurality of datasets; as part of each data ingestion process among the plurality of other data ingestion processes, load a second copy of the user function using a second sandboxed runtime that is associated with each other consumer service; using the second sandboxed runtime that is local to each other consumer service, execute, over records of the datasets, the second copy of the user function directed to the destination table identified in the reference and filter the records or write new records resulting from the function to the destination table.
6. The computer system of claim 1, wherein the hermetic programming language prevents the user function from performing a recursion.
7. The computer system of claim 1, wherein the hermetic programming language prevents the user function from appending to an external list or loading an external script.
8. A computer-implemented method comprising: receiving a digitally defined user function that has been coded using a hermetic and deterministic human-readable programming language, wherein

the hermetic and deterministic programming language prevents the user function from accessing a file system, a network, and a system clock; storing the user function in a centralized database of user functions, in a table of user functions with a reference to a destination table; at a particular consumer service among a plurality of different consumer services executing in a SaaS-based data analytics platform, initiating a data ingestion process; loading a copy of the user function using a sandboxed runtime that is associated with the particular consumer service; using the sandboxed runtime local to the particular consumer service, executing the user function over records directed to the destination table identified in the reference and filtering the records or writing new records resulting from the function to the destination table; the consumer services being further programmed to receive an identifier of a source dataset for ingestion; read successive records of the source dataset until all records have been read; and for each record of the successive records that is read, determine whether a user function is present in the centralized database and has a reference to the destination table, and only when the user function is present and has the reference to the destination table, execute the user function over each record.

9. The method of claim 8, the SaaS-based data analytics platform comprising a customer data platform (CDP).

10. The method of claim 8, further comprising deleting the user function from the particular consumer service or its sandboxed runtime after completing the data ingestion process.

11. The method of claim 8, each of the consumer services being programmed to write a new record resulting from executing the user function only when the new record is valid.

12. The method of claim 8, further comprising, after completing the data ingestion process, deleting the user function from the particular consumer service or its sandboxed runtime.

13. The method of claim 8, further comprising: asynchronously, with respect to the execution of the particular consumer service, using a plurality of other consumer services among the plurality of different consumer services executing in the SaaS-based data analytics platform, initiating a plurality of other data ingestion processes for a plurality of datasets; as part of each data ingestion process among the plurality of other data ingestion processes, loading a second copy of the user function using a second sandboxed runtime that is associated with each other consumer service; using the second sandboxed runtime that is local to each other consumer service, executing, over records of the datasets, the second copy of the user function directed to the destination table identified in the reference and filtering the records or writing new records resulting from the function to the destination table.

14. The method of claim 8, wherein the hermetic programming language prevents the user function from performing a recursion.

15. The method of claim 8, wherein the hermetic programming language prevents the user function from appending to an external list or loading an external script.

16. The method of claim 8, wherein the hermetic programming language prevents the user function from performing more than 500 steps.

17. The method of claim 8, wherein the user function is configured to mask one or more data fields in the records using a hash function executed over the one or more data fields.

18. The method of claim 8, further comprising, after completing the data ingestion process, deleting the user function from the sandboxed runtime.

---