

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393603
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Taylor; Alan L. et al.

Data storage system with efficient data change reporting for replication

Abstract

Selective replication of data in a data storage system includes tracking write transactions using tracking entries in a tracking structure. The tracking structure is periodically processed to create a reporting structure containing reporting entries each identifying a respective extent modified by one or more of the write transactions. The processing for each tracking entry includes (1) examining the reporting structure to determine whether the tracking entry expands an extent of a reporting entry, (2) based on the tracking entry not expanding an extent a reporting entry, creating a new reporting entry based on the tracking entry, and (3) based on the tracking entry expanding an extent of a reporting entry, modifying the reporting entry to identify the extent as so expanded. The reporting structure is used to perform a replication copy operation of the extents identified therein.

Inventors: Taylor; Alan L. (Cary, NC), Seela; Nagapraveen V. (Cary, NC), Jose; Anitta (Portland, OR)

Applicant: Dell Products L.P. (Round Rock, TX)

Family ID: 1000008765097

Assignee: Dell Products L.P. (Round Rock, TX)

Appl. No.: 18/243761

Filed: September 08, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250086198 A1	Mar. 13, 2025

Publication Classification

Int. Cl.: G06F16/27 (20190101); G06F16/23 (20190101); G06F16/25 (20190101)

U.S. Cl.:

CPC **G06F16/27** (20190101); **G06F16/2379** (20190101); **G06F16/252** (20190101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6148368	12/1999	DeKoning	711/E12.019	G06F 12/0866
7647463	12/2009	Jarvis et al.	N/A	N/A
7890469	12/2010	Maionchi	707/654	G06F 16/10
8073922	12/2010	Bates	N/A	N/A
9280591	12/2015	Kharatishvili et al.	N/A	N/A
9633056	12/2016	Attarde et al.	N/A	N/A
2015/0134926	12/2014	Yang	711/167	G06F 3/0652
2016/0306654	12/2015	Diaconu	N/A	G06F 9/4881

Primary Examiner: Zhao; Yu

Attorney, Agent or Firm: BainwoodHuang

Background/Summary

BACKGROUND

(1) The invention is related to the field of data storage systems that provide protective remote replication of data.

SUMMARY

(2) A technique is disclosed for selective replication of data in a data storage system, which includes tracking individual write transactions using respective tracking entries in a tracking structure, each tracking entry identifying an extent of a respective write transaction. The tracking structure is periodically processed to create a reporting structure containing reporting entries each identifying a respective extent modified by one or more of the write transactions. The processing for each tracking entry includes (1) examining the reporting structure to determine whether the tracking entry expands an extent of any of a set of the reporting entries, (2) based on the tracking entry not expanding an extent of any of the set of reporting entries, creating a new reporting entry in the reporting structure based on the tracking entry, and (3) based on the tracking entry expanding an extent of one of the set of reporting entries, modifying the one reporting entry to identify the extent as expanded by the tracking entry. The reporting structure is subsequently used to perform a replication copy operation of the extents identified therein. The processing has a merging characteristic that enhances efficiency of storing and reporting data change information, as well as efficiency of the subsequent copy operation.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The foregoing and other objects, features and advantages will be apparent from the following description of particular embodiments of the invention, as illustrated in the accompanying drawings in which like reference characters refer to the same parts throughout the different views.
- (2) FIG. 1 is a block diagram of a data processing system having data storage systems configured for protective data replication;
- (3) FIG. 2 is a block diagram of a data storage system from a hardware perspective;
- (4) FIG. 3 is a functional block diagram of a data storage system in relevant part;
- (5) FIG. 4 is a schematic illustration of a sequence of replication periods;
- (6) FIG. 5 is a flow diagram of operations of tracking and reporting writes for replication;
- (7) FIG. 6 is a schematic illustration of a sequence of writes occurring in parts of a storage extent;
- (8) FIG. 7 is a schematic depiction of a sequence of states of a reporting structure as the writes of FIG. 6 occur.

DETAILED DESCRIPTION

Overview

- (9) In data storage environments employing asynchronous replication, there may be a Recovery Point Objective (RPO) target for a replication session, which indicates how far behind a destination (replica) is allowed to be from a source (production) volume. For example the RPO may be specified in minutes, and is selected based on the level of protection desired. A data storage system periodically performs replication backups in order to meet the RPO. The source data storage system tracks the changes that need to be replicated to the destination, and the changes are read from the source and written to the destination in each replication cycle. Currently there is a trend toward “Near-Zero” RPO less than 30 seconds.
- (10) One known RPO solution employs snapshot capabilities of a data storage system. Each RPO cycle is begun by creating a snapshot for the next RPO cycle. A copy is performed based on differences between snapshots for the current and previous RPO cycles. A system may use a differencing API to obtain the changes, process those changes to determine the data that actually needs to be copied (replicated), and then replicates the data to the remote destination. However, there can be drawback to the technique of identifying changes based on snapshot differences. For example, expensive flushes of cached data may be required in order to ensure correctness. Because of performance limitations, Near-Zero RPO may not be able to use such a technique for identifying changed data and still maintain a low target RPO (e.g., 30 seconds or less).
- (11) A disclosed technique tracks changed data (writes) using per-core data structures in cache memory that record the volume, offset, and length of each write during a replication cycle. The memory footprint is generally quite small compared to the size of the cache and typical active data set size. The technique employs a reporting structure having a merging characteristic that enhances efficiency of storing and reporting data change information, as well as efficiency of the subsequent copy operation.
- (12) In one embodiment the technique may employ a structure known as a “b-tree”, which resembles a binary search tree and has certain general characteristics. For example, it grows by splitting into two halves and promoting a node, and all its leaf nodes are at the same depth (i.e., it maintains balance). Each node has a certain number of elements, each of which in the present context is an offset and length tuple for a modified extent to be replicated. On insertion of a new element, it is determined whether it can be merged with an existing element, which would be expanded accordingly. For example, adjacent writes are coalesced into a single larger write, and overwrites are tracked as one element, both of these being common occurrences in data storage. In general, the technique reduces the number of elements required to describe a set of changed data

over an RPO cycle, versus a technique that simply records each individual write (offset and length), providing the above-described efficiency benefits.

(13) Other aspects and variations are provided in the detailed description of embodiments below.

Embodiments

(14) FIG. 1 shows a data processing system having a set of host computers (hosts) **10** coupled to data storage systems (DSSs) **12** via a network (NW) **14**. As generally known, the DSSs **12** provide secondary storage services to the hosts **10**, maintaining data storage objects such as volumes, logical disks, etc. that are implemented using underlying physical storage resources (e.g., magnetic media, Flash-programmable memory). The hosts **10** issue data writes and reads to/from these data storage objects, which are processed by the DSSs **12** to store write data and to retrieve stored data and return it in response to read requests.

(15) The system of FIG. 1 employs a technique known as “asynchronous replication” for protection of stored data and ability to return the system to a recent operating condition after a disruption (e.g., failure of a production device). Thus in FIG. 1 the DSSs **12** are identified as a “source”

(production) DSS **12-S** and a “destination” (replica) DSS **12-D**. In operation, a replication session is established between the two DSSs **12** and there is regular copying (replicating) of stored data from source DSS **12-S** to destination DSS **12-D**, as indicated by arrow **16**. Pertinent details of replication operation and the manner of tracking data changes (writes) are described more below.

(16) FIG. 2 is a block diagram of a DSS **12** from a hardware perspective. It includes front-end (FE) interface circuitry **20**, back-end (BE) interface circuitry **22**, and storage processing circuitry **24**. The FE circuitry **20** interfaces the DSS **12** to the network **14** and remote hosts **10** (FIG. 1), and the BE circuitry **22** provides functional connection to local storage devices (DEV) **26** such as magnetic disks, Flash-programmable memory, etc. The storage processing circuitry **24** stores and executes computer program instructions to realize a variety of functional components that provide the functionality of the various data storage services (e.g., primary storage, replication and recovery, data compression and security, etc.), using the other components **20**, **22** and **26** as necessary for this functionality.

(17) FIG. 3 is a functional block diagram of a DSS **12** in relevant part, where the various functions are provided by software-implemented components as described above. In particular, this diagram focuses on data write operations and associated tracking and reporting for replication purposes, as outlined above. Components include a front end **30**, cache **32** and back end **34**, as well as a reporter **36** and copier **38**. The cache **32** employs a tracking structure **40** for tracking host writes **42**. The reporter **36** employs a reporting structure **44** and a scatter-gather list (SGL) generator **46**.

(18) Overall operation is generally as follows, with additional details provided below. The front end **30** receives host writes **42** and applies them to the cache **32** (i.e., stores write data in designated locations as generally known). Under various conditions, data of the cache **32** is “destaged” or written out to backing storage on devices **26**, via the back end **34**. For replication purposes, writes occurring during replication cycles are tracked in the tracking structure **40**, wherein the record or entry for each write includes information including an identifier of the device **26** to which the write is directed, as well as an offset and a length which describe the location of the data modified by the write, as generally known in the art. Also, because of concurrency between tracking and reporting as described below, the write records also include an identifier for the replication cycle in which they occur, which may be realized for example using a snapshot identifier when snapshots are used to bound the replication cycles.

(19) The reporter **36** operates in each replication cycle to process the records of the tracking structure **40** to create records/entries in the reporting structure **44**, which is then used by SGL generator **46** to send change information to the copier **38** in the form of an SGL. The copier **38** uses the information to copy changed data extents from the cache **32** to the destination DSS **12-D**. As generally known, if data has already been destaged to a device **26** then the data is first brought back into the cache **32** and provided to the copier **38** for this replication copying operation.

(20) FIG. 4 is a simplified schematic illustration of ongoing, cyclic replication operations as described above. Operation proceeds in a sequence of cycles or periods **50**, each initiated by a command or operation shown as RP (recovery point). This command is generated by a higher-level controller of replication operation, which is not explicitly shown in the diagrams. At the end of each period **50** certain operations are performed to generate a corresponding report **52** (e.g., report **52-1** for period **50-1**, etc.). The remaining description focuses primarily on details of the reporting, using a simple example to illustrate various structural and functional features.

(21) FIG. 5 is a flow diagram of replication-related operation with particular focus on reporting-related operations such as may be performed by the reporter **36** (FIG. 3).

(22) At **60**, individual write transactions are tracked using respective tracking entries in a tracking structure, wherein each tracking entry identifies an extent of a respective write transaction. In the embodiment of FIG. 3, this tracking is performed by cache **32** using the tracking structure **40**. The term “extent” refers to the data that is written, identified with an offset (starting address) and length (e.g. in bytes). As noted above, each tracking entry may also include additional information such as the device ID and period ID (e.g., snapshot ID for the replication cycle/period).

(23) At **62**, the tracking structure is periodically processed to create a reporting structure, which is populated with reporting entries each identifying a respective extent modified by one or more of the write transactions. The processing for each tracking entry includes (1) examining the reporting structure to determine whether the tracking entry expands an extent of any of a set of the reporting entries (i.e., enlarge the size of such an extent), (2) based on the tracking entry not expanding an extent of any of the set of reporting entries, creating a new reporting entry in the reporting structure based on the tracking entry, and (3) based on the tracking entry expanding an extent of one of the set of reporting entries, modifying the one reporting entry to identify the extent as expanded by the tracking entry. It will be understood that this modifying operation, which may also be referred to as “merging” herein, provides part of the above-described efficiencies of storage and copying. In the embodiment of FIG. 3, this processing is performed by the reporter **36** using the reporting structure **44**. All of these operations and attendant benefits are further described with reference to an example below.

(24) At **64**, the reporting structure is used to perform a replication copy operation of the extents identified therein. That is, the extents as represented in the reporting structure are the subject of the replication copying for that period. In the embodiment of FIG. 3, these operations are performed jointly by the reporter **36** and copier **38**. The SGL generator **46** of reporter **36** uses the reporting entries of reporting structure **44** to create the SGL and provide it to the copier **38**, which then performs the copy operation for these extents by reading the data from the cache **32** and sending it to the destination DSS **12-D** for writing to a replica device.

(25) FIGS. 6 and 7 describe a simplified example used to illustrate the above operations.

(26) FIG. 6 shows an area **70** of storage of a device **26** and a pattern of five writes **72-1** through **72-5** performed in sequence. Writes **72-1** through **72-3** are directed to respective disjoint areas. Write **72-4** is to a range completely enveloping the range of write **72-2**, and write **72-5** is to a range that overlaps part of the range of write **72-4** and completely envelopes the range of write **72-3**. As indicated at left, each write is represented by an Offset (O) and Length (L) tuple [O_x; L_x]. As these writes **72** occur, they are logged individually in the tracking structure **40**.

(27) On the right side of FIG. 6 is shown the manner in which these five writes are eventually reported. The range for write **72-1** is reported as a first extent **74** with an offset-length of O₁; L₁, meaning the offset value from write **72-1** and the length value from write **72-1** (i.e., this range is identical to the range written by write **72-1**). Because the result of writes **72-2** through **72-5** is one large contiguous region extending across all the regions for these writes, they are represented by a single second extent **74** that has offset value O₄ (the offset of write **72-4**) and a length value shown as “L>”, which means it is an aggregated length that is computed based on the offsets and lengths of its constituent writes. In this example, omitting intermediate steps (which are described below),

the value $L >$ is equal to $(O5 - O4) + L5$, where the term $(O5 - O4)$ captures the portion of write **72-4** that does not overlap with write **72-5**.

(28) FIG. 7 illustrates the contents of the reporting structure **44** as the writes **72-1** through **72-5** occur in this example. In this case the reporting structure **44** is realized using a so-called “b-tree”, which is a generalization of a binary search tree. B-trees have various properties that make them suitable for this kind of application, i.e., capturing a pattern of data items scattered throughout a region. As explained below, standard b-tree manipulation is augmented to provide desired expanding (or merging) of existing entries to increase efficiency especially of the ensuing copy operation. In FIG. 7, the b-tree itself is indicated by reference **80**, and its individual entries by the letter E. For each write **72-1** through **72-5** there is a corresponding resulting state of the b-tree, **80-1** through **80-5**. Each entry E has an entry number (1, 2, . . .) generally reflecting the sequence in which they are created (E1, E2, . . .). Also shown for writes **72-3** through **72-5** are corresponding notional intermediate states of the b-tree **80**, identified as **82-3** through **82-5**. Each entry E has an offset-length tuple $[Ox; Ly]$ following the scheme of FIG. 6, i.e., x and y are taken from corresponding ones of the writes **72-1** through **72-5** or are calculated, which is represented by “>”. In this simplified example, it is assumed that the tree **80** is binary and that each node can store up to two entries Ex.

(29) For write **72-1**, a first entry E1 is created with offset-length of $[O1; L1]$ and stored in the tree **80**, resulting in the tree state **80-1** having a single node storing E1.

(30) Once the tree **80** is populated with at least one node, then in subsequent insertions there is an initial step of examining certain nodes to determine if they can be expanded to incorporate some or all of the range of a subsequent write being processed. As noted, this may also be referred to as “merging.” Thus for write **72-2**, for example, its values are compared with those of the single entry E1. In this case, the ranges are disjoint, so no expanding is possible. Thus, a second entry E2 is created with offset-length of $[O2; L2]$ and stored in the tree **80**, resulting in tree state **80-2**. The entry E2 is added to the one node already storing entry E1. E2 is added to the right of E1, due to its location in a higher address range.

(31) For write **72-3**, there is a notional intermediate state **82-3** in which entry E3 is added to the right of entry E2, due to its higher address range (note that the range for write **72-3** does not adjoin or overlap either E1 or E2, so no expanding/merging is possible). However, because the single node is already full, a new node must be added. The existing node is logically “split” into two same-level nodes, one storing E1 and the other storing E3. The notionally intermediate entry E2 is elevated one level and its node is given pointers to the nodes storing E1 and E3 (tree state **80-3**).

(32) At this point no efficiency has been gained, as there are three entries E1-E3 for the three writes **72-1** through **72-3**, due to the lack of expanding/merging opportunities. However, this changes as writes **72-4** and **72-5** occur, as explained below.

(33) For write **82-4**, there is a notionally intermediate state **82-4** in which E4 fits between E1 and E2 due to the position of its range (it could be part of the E1 node as shown or of the E2 node). However, as shown in FIG. 6, write **72-4** completely envelopes the range for write **72-2**, which is represented in entry E2. Thus, entry E2 can be expanded to represent the entire range covered by writes **72-2** and **72-4**, which in this case is the same as the whole range of write **72-4**. The result is tree state **80-4** in which entry E2 now has values $[O4; L4]$.

(34) The processing is similar for write **72-5** which has notional intermediate state **82-5**. In this case multiple merges occur. Referring again to FIG. 6, write **72-5** partially overlaps write **72-4** (which is represented by existing entry E2) and completely envelopes write **72-3** (represented by entry E3). The merging may occur in a variety of ways depending on implementation. In one example approach, it may occur as follows: 1. Write **72-5** merged into entry E3 (E3 expanded), based on overlap. 2. Entry E3 merged into entry E2 (E2 expanded), based on overlap. E2 has offset $O4$ and the calculated length $L >$ as described above with reference to FIG. 6. 3. With the tree **80-5** reduced to only two entries E1, E2, it can be collapsed to one level holding only one 2-entry node,

which is depicted in FIG. 7 (tree state **80-5**). Alternatively, the two-level structure may be maintained, with the right-hand pointer for the root node (storing E2) being a null value (due to deletion of entry E3 after merging).

(35) The above-mentioned efficiencies will be apparent based on the above example. The reporting structure **44** in the form of b-tree **80** finally stores only two entries E1, E2 representing the two extents **74**, **76** that result from five writes **72-1** through **72-5**. Only two copying operations are involved rather than five, and there is no duplicated copying of sub-regions corresponding to overlaps (e.g., the entire region of writes **72-2** and **72-3**, and the overlap between writes **72-4** and **72-5**).

(36) The above example also illustrates that merging occurs both horizontally (nodes at same tree level) as well as vertically (nodes at different levels), as described above for write **72-5** for example.

(37) The above uses a conventional approach of splitting nodes “in the middle”, i.e., dividing their entries equally between two new nodes. Because the entries Ex represent ranges, there is a notion of “distance” between the entries and thus their nodes, which could be used to perform splitting in some different way that might provide an operational benefit. If highly sequential data is detected then there could be a decision to promote not the middle element but rather an element positioned mid-way in terms of its range relative to ranges of other elements. This may provide for nodes to remain fuller over time and make more efficient use of storage space for the reporting structure.

(38) While various embodiments of the invention have been particularly shown and described, it will be understood by those skilled in the art that various changes in form and details may be made therein without departing from the scope of the invention as defined by the appended claims.

Claims

1. A method of selective replication of host data stored in a data storage system, comprising: establishing and maintaining a replication session with a second data storage system; and as part of the replication session, in each of ongoing replication cycles in which write transactions store newly written host data in the data storage system, (1) creating a reporting structure identifying the newly written host data to be replicated, and (2) copying the newly written host data identified in the reporting structure to the second data storage system for replication, wherein creating the reporting structure includes: (i) tracking individual write transactions using respective tracking entries (TEs) in a tracking structure, each tracking entry identifying a respective TE extent of a respective write transaction; and (ii) processing the tracking structure to create the reporting structure containing reporting entries (REs) each identifying a respective RE extent modified by one or more of the write transactions, the processing for each tracking entry including (a) examining the reporting structure to determine whether the TE extent of the tracking entry expands, by overlapping, an RE extent of any of a set of the reporting entries, (b) based on the TE extent of the tracking entry not expanding an RE extent of any of the set of reporting entries, creating a new reporting entry in the reporting structure based on the tracking entry, and (c) based on the TE extent of the tracking entry expanding an existing RE extent of one of the set of reporting entries, modifying the one reporting entry to identify an expanded RE extent spanning the existing RE extent and the TE extent of the tracking entry, the modifying of the one reporting entry reducing an amount of the newly written host data copied in the copying step by an amount of overlap between the existing RE extent and the TE extent of the tracking entry.
2. The method of claim 1, wherein each tracking entry includes an offset and length for the TE extent of the write transaction, and further includes a device identifier and replication period identifier, the device identifier identifying a storage device to which the write transaction is directed, the replication period identifier identifying a specific replication period in which the write occurred and for which the tracking, reporting and copying operations are performed.

3. The method of claim 2, wherein the replication period identifier is a snapshot identifier for a data snapshot taken at a boundary of the replication period.
4. The method of claim 1, wherein the copying includes generating a scatter-gather list from the reporting entries of the reporting structure and providing the scatter-gather list to a copier component for copying RE data extents identified therein.
5. The method of claim 1, wherein determining whether the tracking entry expands an RE extent of any of the set of the reporting entries includes determining whether the tracking entry is directed to an extent either adjoining or overlapping an RE extent of any of the set of reporting entries.
6. The method of claim 1, wherein an expanded reporting entry is created based on a plurality of constituent write transactions that have modified the TE extents thereof, and includes (1) an offset taken from a lowest-range one of the constituent write transactions, and (2) an aggregated length value calculated using respective offsets and lengths of the constituent write transactions.
7. The method of claim 1, wherein the reporting structure includes a b-tree employing a plurality of tree-organized, multi-entry nodes in which entries are ordered according to locations of their respective RE extents, and wherein creating a new reporting entry includes one of (1) adding the new reporting entry to an existing node, or (2) creating a new node and adding the new reporting entry to the new node.
8. The method of claim 7, wherein creating a new node occurs as part of a node splitting operation for an existing node, the node splitting resulting in (1) the new node as a peer of the existing node, and (2) a new next-level node pointing to both the existing node and the new nodes, and wherein the new reporting entry and existing reporting entries of the existing node are distributed among the existing node, new node and new next-level node in a predetermined manner.
9. The method of claim 8, wherein (1) the new next-level node receives a middle reporting entry among the reporting entries being distributed, (2) the existing node receives one or more of the existing reporting entries located below the middle reporting entry, and (3) the new node receives the new reporting entry.
10. The method of claim 7, wherein modifying the one reporting entry includes expanding its RE extent to incorporate part or all of an RE extent of a peer reporting entry at a same level of the b-tree.
11. The method of claim 7, wherein modifying the one reporting entry includes expanding its RE extent to incorporate part or all of an RE extent of a child reporting entry at a next lower level of the b-tree.
12. A data storage system, comprising: storage devices for storing host data, front-end (FE) circuitry providing an interface to a network and remote hosts; back-end (BE) circuitry providing functional connection to the storage devices; and storage processing circuitry storing and executing computer program instructions to cause the data storage system to selectively replicate stored host data to a remote replica system, by: establishing and maintaining a replication session with the remote replica system; and as part of the replication session, in each of ongoing replication cycles in which host write transactions store newly written host data in the data storage devices, (1) creating a reporting structure identifying the newly written host data to be replicated, and (2) copying the newly written host data identified in the reporting structure to the remote replica system for replication, wherein creating the reporting structure includes: (i) tracking individual write transactions using respective tracking entries (TEs) in a tracking structure, each tracking entry identifying a respective TE extent of a respective write transaction; and (ii) processing the tracking structure to create the reporting structure containing reporting entries (REs) each identifying a respective RE extent modified by one or more of the write transactions, the processing for each tracking entry including (a) examining the reporting structure to determine whether the TE extent of the tracking entry expands, by overlapping, an RE extent of any of a set of the reporting entries, (b) based on the TE extent of the tracking entry not expanding an RE extent of any of the set of reporting entries, creating a new reporting entry in the reporting structure based on the tracking

entry, and (c) based on the TE extent of the tracking entry expanding an existing RE extent of one of the set of reporting entries, modifying the one reporting entry to identify an expanded RE extent spanning the existing RE extent and the TE extent of the tracking entry, the modifying of the one reporting entry reducing an amount of the newly written host data copied in the copying step by an amount of overlap between the existing RE extent and the TE extent of the tracking entry.

13. The data storage system of claim 12, wherein each tracking entry includes an offset and length for the TE extent of the write transaction, and further includes a device identifier and replication period identifier, the device identifier identifying a storage device to which the write transaction is directed, the replication period identifier identifying a specific replication period in which the write occurred and for which the tracking, reporting and copying operations are performed.

14. The data storage system of claim 12, wherein determining whether the tracking entry expands an RE extent of any of the set of the reporting entries includes determining whether the tracking entry is directed to an extent either adjoining or overlapping an RE extent of any of the set of reporting entries.

15. The data storage system of claim 12, wherein an expanded reporting entry is created based on a plurality of constituent write transactions that have modified the TE extents thereof, and includes (1) an offset taken from a lowest-range one of the constituent write transactions, and (2) an aggregated length value calculated using respective offsets and lengths of the constituent write transactions.

16. The data storage system of claim 12, wherein the reporting structure includes a b-tree employing a plurality of tree-organized, multi-entry nodes in which entries are ordered according to locations of their respective RE extents, and wherein creating a new reporting entry includes one of (1) adding the new reporting entry to an existing node, or (2) creating a new node and adding the new reporting entry to the new node.

17. The data storage system of claim 16, wherein creating a new node occurs as part of a node splitting operation for an existing node, the node splitting resulting in (1) the new node as a peer of the existing node, and (2) a new next-level node pointing to both the existing node and the new nodes, and wherein the new reporting entry and existing reporting entries of the existing node are distributed among the existing node, new node and new next-level node in a predetermined manner.

18. The data storage system of claim 17, wherein (1) the new next-level node receives a middle reporting entry among the reporting entries being distributed, (2) the existing node receives one or more of the existing reporting entries located below the middle reporting entry, and (3) the new node receives the new reporting entry.

19. The data storage system of claim 16, wherein modifying the one reporting entry includes expanding its RE extent to incorporate part or all of an RE extent of a peer reporting entry at a same level of the b-tree.

20. The data storage system of claim 16, wherein modifying the one reporting entry includes expanding its RE extent to incorporate part or all of an RE extent of a child reporting entry at a next lower level of the b-tree.
