

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393608
Kind Code	B1
Date of Patent	August 19, 2025
Inventor(s)	Ma; Jiyue

Metastore manager for multi-platform data operations

Abstract

A metastore manager facilitates access to data/metadata stored on at least one data management platform. A first data processing operation is initiated, in association with a data set maintained in a first data store associated with a first data management platform, to generate a first processed data set. The first processed data set is stored in the first data store and a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store. The first metadata set includes partition metadata indicative of partitioning information associated with the first processed data set. A synchronization operation is implemented to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set. The metastore manager may facilitate access, by a data processing pipeline, to the processed data set.

Inventors:	Ma; Jiyue (Hangzhou, CN)
Applicant:	Zoom Communications, Inc. (San Jose, CA)
Family ID:	1000007953823
Assignee:	Zoom Communications, Inc. (San Jose, CA)
Appl. No.:	18/673931
Filed:	May 24, 2024

Publication Classification

Int. Cl.: G06F7/00 (20060101); G06F16/242 (20190101); G06F16/27 (20190101)

U.S. Cl.:

CPC G06F16/278 (20190101); G06F16/2433 (20190101);

Field of Classification Search

CPC: G06F (16/278); G06F (16/2433); G06F (16/258); G06F (16/119); G06F (21/604); G06F (16/219); G06F (21/6245); G06F (21/6227)

USPC: 707/602; 707/689; 707/713; 707/737; 707/769; 707/802; 707/999.1; 707/17.005; 707/610; 707/624; 707/694; 707/770; 707/999.101; 711/100; 711/113; 711/154; 711/162

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8200613	12/2011	Wan et al.	N/A	N/A
10102269	12/2017	Marin	N/A	N/A
10120904	12/2017	Ranganathan	N/A	N/A
10541938	12/2019	Timmerman et al.	N/A	N/A
11074107	12/2020	Nandakumar	N/A	N/A
11107177	12/2020	Ashley	N/A	G06T 1/20
11153173	12/2020	Rebeja	N/A	N/A
11741119	12/2022	Orun	N/A	N/A
11941014	12/2023	Das	N/A	G06F 16/219
2011/0282835	12/2010	Cannon	707/622	G06F 16/119
2018/0196867	12/2017	Wiesmaier et al.	N/A	N/A
2020/0257596	12/2019	Gokhale	N/A	G06F 16/25
2021/0124739	12/2020	Karanasos et al.	N/A	N/A
2021/0294814	12/2020	Murthy	N/A	N/A
2022/0138004	12/2021	Nandakumar	N/A	N/A
2024/0427911	12/2023	Apostolatos	N/A	G06F 16/258

Primary Examiner: Uddin; Md I

Attorney, Agent or Firm: Young Basile Hanlon & MacFarlane, P.C.

Background/Summary

FIELD

(1) This disclosure generally relates to data management and, more specifically, to metastore managers for multi-platform data operations.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) This disclosure is best understood from the following detailed description when read in conjunction with the accompanying drawings. It is emphasized that, according to common practice, the various features of the drawings are not to-scale. On the contrary, the dimensions of the various features are arbitrarily expanded or reduced for clarity.
- (2) FIG. 1 is a block diagram of an example of an electronic computing and communications system.
- (3) FIG. 2 is a block diagram of an example internal configuration of a computing device of an

electronic computing and communications system.

(4) FIG. 3 is a block diagram of an example of a software platform implemented by an electronic computing and communications system.

(5) FIG. 4 is a block diagram of an example of a Spark execution system for executing structured query language (SQL) queries.

(6) FIG. 5 is an example of a Spark optimization framework for translating one or more SQL queries of SQL code into executable operations on a database.

(7) FIG. 6 is a diagram illustrating an example of a data management environment.

(8) FIG. 7 is a block diagram of an example of a metastore manager system.

(9) FIG. 8 is a schematic block diagram illustrating an example of an implementation of a multi-platform data management environment in accordance with the present disclosure.

(10) FIG. 9 is a schematic block diagram illustrating another example of an implementation of a multi-platform data management environment in accordance with the present disclosure.

(11) FIG. 10 is a flowchart of an example of a technique for performing a multi-platform data operation.

DETAILED DESCRIPTION

(12) Conferencing software is frequently used across various industries to support video-enabled conferences between participants in multiple locations. In some cases, each of the conference participants separately connects to the conferencing software from their own remote locations. In other cases, one or more of the conference participants may be physically located in and connect to the conferencing software from a conference room or similar physical space (e.g., in an office setting) while other conference participants connect to the conferencing software from one or more remote locations. Conferencing software thus enables people to conduct video conferences without requiring them to be physically present with one another. Conferencing software may be available as a standalone software product or it may be integrated within a software platform, such as a unified communications as a service (UCaaS) platform.

(13) To support conferencing services, a UCaaS platform may be implemented in connection with one or more databases configured to maintain data. Data may include, for example, customer data (e.g., data provided by, about, or otherwise associated with one or more customers), communication data (e.g., data associated with prior conferences, networks, devices, and/or other aspects of implementations of the UCaaS platform. Any number of different data management platforms may be used for ingesting, processing, storing, and/or otherwise managing data. For example, cloud services may be used for data management. In some implementations, several different data management platforms may be utilized. In some cases, a data processing operation may be configured to be executed on one platform in connection with data that is stored in another platform. In some cases, modifying data in one platform may require modifying related data stored on another platform. Because the different data management platforms may be provided by different entities and/or services, the different data management platforms may not be structured similarly. As a result, establishing dependencies between data management platforms to facilitate this sort of cross-platform management can be challenging and result in data loss and/or corruption.

(14) Implementations of this disclosure address problems such as these by providing a metastore manager in association with at least one data management platform. The metastore manager may be configured to facilitate access to data and/or metadata stored on one or more data management platforms. The access to the data may be access by users and/or by other platforms. For example, the metastore manager may facilitate access, by a first data management platform, to data and/or metadata stored on a second data management platform. In some cases, the metastore manager may facilitate access, by a data processing pipeline, to data stored on one or more of the data management platform. In this way, aspects of the disclosure may facilitate cross-platform data management operations, enabling independent platforms to exchange data.

(15) In some examples of the present disclosure, implementations may include or otherwise use

one or more artificial intelligence or machine learning (collectively, AI/ML) systems having one or more models trained for one or more purposes. Use or inclusion of such AI/ML systems, such as for implementation of certain features or functions, may be turned off by default, where a user, an organization, or both must opt-in to utilize the features or functions that include or otherwise use an AI/ML system. User or organizational consent to use the AI/ML systems or features may be provided in one or more ways, for example, as explicit permission granted by a user prior to using an AI/ML feature, as administrative consent configured by administrator settings, or both. Users for whom such consent is obtained can be notified that they will be interacting with one or more AI/ML systems or features, for example, by an electronic message (e.g., delivered via a chat or email service or presented within a client application or webpage) or by an on-screen prompt, which can be applied on a per-interaction basis. Those users can also be provided with an easy way to withdraw their user consent, for example, using a form or like element provided within a client application, webpage, or on-screen prompt to allow individual users to opt-out of use of the AI/ML systems or features.

(16) To enhance privacy and safety, as well as provide other benefits, the AI/ML processing system may be prevented from using a user's or organization's personal information (e.g., audio, video, chat, screen-sharing, attachments, or other communications-like content (such as poll results, whiteboards, or reactions)) to train any AI/ML models and instead only use the personal information for inference operations of the AI/ML processing system. Instead of using the personal information to train AI/ML models, AI/ML models may be trained using one or more commercially licensed data sets that do not contain the personal information of the user or organization.

(17) To describe some implementations in greater detail, reference is first made to examples of hardware and software structures used to implement a system for multi-platform data operations. FIG. 1 is a block diagram of an example of an electronic computing and communications system **100**, which can be or include a distributed computing system (e.g., a client-server computing system), a cloud computing system, a clustered computing system, or the like.

(18) The system **100** includes one or more customers, such as customers **102A** through **102B**, which may each be a public entity, private entity, or another corporate entity or individual that purchases or otherwise uses software services, such as of a UCaaS platform provider. Each customer can include one or more clients. For example, as shown and without limitation, the customer **102A** can include clients **104A** through **104B**, and the customer **102B** can include clients **104C** through **104D**. A customer can include a customer network or domain. For example, and without limitation, the clients **104A** through **104B** can be associated or communicate with a customer network or domain for the customer **102A** and the clients **104C** through **104D** can be associated or communicate with a customer network or domain for the customer **102B**.

(19) A client, such as one of the clients **104A** through **104D**, may be or otherwise refer to one or both of a client device or a client application. Where a client is or refers to a client device, the client can comprise a computing system, which can include one or more computing devices, such as a mobile phone, a tablet computer, a laptop computer, a notebook computer, a desktop computer, or another suitable computing device or combination of computing devices. Where a client instead is or refers to a client application, the client can be an instance of software running on a customer device (e.g., a client device or another device). In some implementations, a client can be implemented as a single physical unit or as a combination of physical units. In some implementations, a single physical unit can include multiple clients.

(20) The system **100** can include a number of customers and/or clients or can have a configuration of customers or clients different from that generally illustrated in FIG. 1. For example, and without limitation, the system **100** can include hundreds or thousands of customers, and at least some of the customers can include or be associated with a number of clients.

(21) The system **100** includes a datacenter **106**, which may include one or more servers. The datacenter **106** can represent a geographic location, which can include a facility, where the one or

more servers are located. The system **100** can include a number of datacenters and servers or can include a configuration of datacenters and servers different from that generally illustrated in FIG. **1**. For example, and without limitation, the system **100** can include tens of datacenters, and at least some of the datacenters can include hundreds or another suitable number of servers. In some implementations, the datacenter **106** can be associated or communicate with one or more datacenter networks or domains, which can include domains other than the customer domains for the customers **102A** through **102B**.

(22) The datacenter **106** includes servers used for implementing software services of a UCaaS platform. The datacenter **106** as generally illustrated includes an application server **108**, a database server **110**, and a telephony server **112**. The servers **108** through **112** can each be a computing system, which can include one or more computing devices, such as a desktop computer, a server computer, or another computer capable of operating as a server, or a combination thereof. A suitable number of each of the servers **108** through **112** can be implemented at the datacenter **106**. The UCaaS platform uses a multi-tenant architecture in which installations or instantiations of the servers **108** through **112** is shared amongst the customers **102A** through **102B**.

(23) In some implementations, one or more of the servers **108** through **112** can be a non-hardware server implemented on a physical device, such as a hardware server. In some implementations, a combination of two or more of the application server **108**, the database server **110**, and the telephony server **112** can be implemented as a single hardware server or as a single non-hardware server implemented on a single hardware server. In some implementations, the datacenter **106** can include servers other than or in addition to the servers **108** through **112**, for example, a media server, a proxy server, or a web server.

(24) The application server **108** runs web-based software services deliverable to a client, such as one of the clients **104A** through **104D**. As described above, the software services may be of a UCaaS platform. For example, the application server **108** can implement all or a portion of a UCaaS platform, including conferencing software, messaging software, and/or other intra-party or inter-party communications software. The application server **108** may, for example, be or include a unitary Java Virtual Machine (JVM).

(25) In some implementations, the application server **108** can include an application node, which can be a process executed on the application server **108**. For example, and without limitation, the application node can be executed in order to deliver software services to a client, such as one of the clients **104A** through **104D**, as part of a software application. The application node can be implemented using processing threads, virtual machine instantiations, or other computing features of the application server **108**. In some such implementations, the application server **108** can include a suitable number of application nodes, depending upon a system load or other characteristics associated with the application server **108**. For example, and without limitation, the application server **108** can include two or more nodes forming a node cluster. In some such implementations, the application nodes implemented on a single application server **108** can run on different hardware servers.

(26) The database server **110** stores, manages, or otherwise provides data for delivering software services of the application server **108** to a client, such as one of the clients **104A** through **104D**. In particular, the database server **110** may implement one or more databases, tables, or other information sources suitable for use with a software application implemented using the application server **108**. The database server **110** may include a data storage unit accessible by software executed on the application server **108**. A database implemented by the database server **110** may be a relational database management system (RDBMS), an object database, an XML database, a configuration management database (CMDB), a management information base (MIB), one or more flat files, other suitable non-transient storage mechanisms, or a combination thereof. The system **100** can include one or more database servers, in which each database server can include one, two, three, or another suitable number of databases configured as or comprising a suitable database type

or combination thereof.

(27) In some implementations, one or more databases, tables, other suitable information sources, or portions or combinations thereof may be stored, managed, or otherwise provided by one or more of the elements of the system **100** other than the database server **110**, for example, one of the clients **104A** through **104D** or the application server **108**.

(28) The telephony server **112** enables network-based telephony and web communications from and/or to clients of a customer, such as the clients **104A** through **104B** for the customer **102A** or the clients **104C** through **104D** for the customer **102B**. For example, one or more of the clients **104A** through **104D** may be voice over internet protocol (VOIP)-enabled devices configured to send and receive calls over a network **114**. The telephony server **112** includes a session initiation protocol (SIP) zone and a web zone. The SIP zone enables a client of a customer, such as the customer **102A** or **102B**, to send and receive calls over the network **114** using SIP requests and responses. The web zone integrates telephony data with the application server **108** to enable telephony-based traffic access to software services run by the application server **108**. Given the combined functionality of the SIP zone and the web zone, the telephony server **112** may be or include a cloud-based private branch exchange (PBX) system.

(29) The SIP zone receives telephony traffic from a client of a customer and directs same to a destination device. The SIP zone may include one or more call switches for routing the telephony traffic. For example, to route a VOIP call from a first VOIP-enabled client of a customer to a second VOIP-enabled client of the same customer, the telephony server **112** may initiate a SIP transaction between a first client and the second client using a PBX for the customer. However, in another example, to route a VOIP call from a VOIP-enabled client of a customer to a client or non-client device (e.g., a desktop phone which is not configured for VOIP communication) which is not VOIP-enabled, the telephony server **112** may initiate a SIP transaction via a VOIP gateway that transmits the SIP signal to a public switched telephone network (PSTN) system for outbound communication to the non-VOIP-enabled client or non-client phone. Hence, the telephony server **112** may include a PSTN system and may in some cases access an external PSTN system.

(30) The telephony server **112** includes one or more session border controllers (SBCs) for interfacing the SIP zone with one or more aspects external to the telephony server **112**. In particular, an SBC can act as an intermediary to transmit and receive SIP requests and responses between clients or non-client devices of a given customer with clients or non-client devices external to that customer. When incoming telephony traffic for delivery to a client of a customer, such as one of the clients **104A** through **104D**, originating from outside the telephony server **112** is received, a SBC receives the traffic and forwards it to a call switch for routing to the client.

(31) In some implementations, the telephony server **112**, via the SIP zone, may enable one or more forms of peering to a carrier or customer premise. For example, Internet peering to a customer premise may be enabled to ease the migration of the customer from a legacy provider to a service provider operating the telephony server **112**. In another example, private peering to a customer premise may be enabled to leverage a private connection terminating at one end at the telephony server **112** and at the other end at a computing aspect of the customer environment. In yet another example, carrier peering may be enabled to leverage a connection of a peered carrier to the telephony server **112**.

(32) In some such implementations, a SBC or telephony gateway within the customer environment may operate as an intermediary between the SBC of the telephony server **112** and a PSTN for a peered carrier. When an external SBC is first registered with the telephony server **112**, a call from a client can be routed through the SBC to a load balancer of the SIP zone, which directs the traffic to a call switch of the telephony server **112**. Thereafter, the SBC may be configured to communicate directly with the call switch.

(33) The web zone receives telephony traffic from a client of a customer, via the SIP zone, and directs same to the application server **108** via one or more Domain Name System (DNS)

resolutions. For example, a first DNS within the web zone may process a request received via the SIP zone and then deliver the processed request to a web service which connects to a second DNS at or otherwise associated with the application server **108**. Once the second DNS resolves the request, it is delivered to the destination service at the application server **108**. The web zone may also include a database for authenticating access to a software application for telephony traffic processed within the SIP zone, for example, a softphone.

(34) The clients **104A** through **104D** communicate with the servers **108** through **112** of the datacenter **106** via the network **114**. The network **114** can be or include, for example, the Internet, a local area network (LAN), a wide area network (WAN), a virtual private network (VPN), or another public or private means of electronic computer communication capable of transferring data between a client and one or more servers. In some implementations, a client can connect to the network **114** via a communal connection point, link, or path, or using a distinct connection point, link, or path. For example, a connection point, link, or path can be wired, wireless, use other communications technologies, or a combination thereof.

(35) The network **114**, the datacenter **106**, or another element, or combination of elements, of the system **100** can include network hardware such as routers, switches, other network devices, or combinations thereof. For example, the datacenter **106** can include a load balancer **116** for routing traffic from the network **114** to various servers associated with the datacenter **106**. The load balancer **116** can route, or direct, computing communications traffic, such as signals or messages, to respective elements of the datacenter **106**.

(36) For example, the load balancer **116** can operate as a proxy, or reverse proxy, for a service, such as a service provided to one or more remote clients, such as one or more of the clients **104A** through **104D**, by the application server **108**, the telephony server **112**, and/or another server. Routing functions of the load balancer **116** can be configured directly or via a DNS. The load balancer **116** can coordinate requests from remote clients and can simplify client access by masking the internal configuration of the datacenter **106** from the remote clients.

(37) In some implementations, the load balancer **116** can operate as a firewall, allowing or preventing communications based on configuration settings. Although the load balancer **116** is depicted in FIG. 1 as being within the datacenter **106**, in some implementations, the load balancer **116** can instead be located outside of the datacenter **106**, for example, when providing global routing for multiple datacenters. In some implementations, load balancers can be included both within and outside of the datacenter **106**. In some implementations, the load balancer **116** can be omitted.

(38) FIG. 2 is a block diagram of an example internal configuration of a computing device **200** of an electronic computing and communications system. In one configuration, the computing device **200** may implement one or more of the client **104**, the application server **108**, the database server **110**, or the telephony server **112** of the system **100** shown in FIG. 1.

(39) The computing device **200** includes components or units, such as a processor **202**, a memory **204**, a bus **206**, a power source **208**, peripherals **210**, a user interface **212**, a network interface **214**, other suitable components, or a combination thereof. One or more of the memory **204**, the power source **208**, the peripherals **210**, the user interface **212**, or the network interface **214** can communicate with the processor **202** via the bus **206**.

(40) The processor **202** is a central processing unit, such as a microprocessor, and can include single or multiple processors having single or multiple processing cores. Alternatively, the processor **202** can include another type of device, or multiple devices, configured for manipulating or processing information. For example, the processor **202** can include multiple processors interconnected in one or more manners, including hardwired or networked. The operations of the processor **202** can be distributed across multiple devices or units that can be coupled directly or across a local area or other suitable type of network. The processor **202** can include a cache, or cache memory, for local storage of operating data or instructions.

(41) The memory **204** includes one or more memory components, which may each be volatile memory or non-volatile memory. For example, the volatile memory can be random access memory (RAM) (e.g., a DRAM module, such as DDR SDRAM). In another example, the non-volatile memory of the memory **204** can be a disk drive, a solid state drive, flash memory, or phase-change memory. In some implementations, the memory **204** can be distributed across multiple devices. For example, the memory **204** can include network-based memory or memory in multiple clients or servers performing the operations of those multiple devices.

(42) The memory **204** can include data for immediate access by the processor **202**. For example, the memory **204** can include executable instructions **216**, application data **218**, and an operating system **220**. The executable instructions **216** can include one or more application programs, which can be loaded or copied, in whole or in part, from non-volatile memory to volatile memory to be executed by the processor **202**. For example, the executable instructions **216** can include instructions for performing some or all of the techniques of this disclosure. The application data **218** can include user data, database data (e.g., database catalogs or dictionaries), or the like. In some implementations, the application data **218** can include functional programs, such as a web browser, a web server, a database server, another program, or a combination thereof. The operating system **220** can be, for example, Microsoft Windows®, Mac OS X®, or Linux®; an operating system for a mobile device, such as a smartphone or tablet device; or an operating system for a non-mobile device, such as a mainframe computer.

(43) The power source **208** provides power to the computing device **200**. For example, the power source **208** can be an interface to an external power distribution system. In another example, the power source **208** can be a battery, such as where the computing device **200** is a mobile device or is otherwise configured to operate independently of an external power distribution system. In some implementations, the computing device **200** may include or otherwise use multiple power sources. In some such implementations, the power source **208** can be a backup battery.

(44) The peripherals **210** includes one or more sensors, detectors, or other devices configured for monitoring the computing device **200** or the environment around the computing device **200**. For example, the peripherals **210** can include a geolocation component, such as a global positioning system location unit. In another example, the peripherals can include a temperature sensor for measuring temperatures of components of the computing device **200**, such as the processor **202**. In some implementations, the computing device **200** can omit the peripherals **210**.

(45) The user interface **212** includes one or more input interfaces and/or output interfaces. An input interface may, for example, be a positional input device, such as a mouse, touchpad, touchscreen, or the like; a keyboard; or another suitable human or machine interface device. An output interface may, for example, be a display, such as a liquid crystal display, a cathode-ray tube, a light emitting diode display, or other suitable display.

(46) The network interface **214** provides a connection or link to a network (e.g., the network **114** shown in FIG. 1). The network interface **214** can be a wired network interface or a wireless network interface. The computing device **200** can communicate with other devices via the network interface **214** using one or more network protocols, such as using Ethernet, transmission control protocol (TCP), internet protocol (IP), power line communication, an IEEE 802.X protocol (e.g., Wi-Fi, Bluetooth, or ZigBee), infrared, visible light, general packet radio service (GPRS), global system for mobile communications (GSM), code-division multiple access (CDMA), Z-Wave, another protocol, or a combination thereof.

(47) FIG. 3 is a block diagram of an example of a software platform **300** implemented by an electronic computing and communications system, for example, the system **100** shown in FIG. 1. The software platform **300** is a UCaaS platform accessible by clients of a customer of a UCaaS platform provider, for example, the clients **104A** through **104B** of the customer **102A** or the clients **104C** through **104D** of the customer **102B** shown in FIG. 1. The software platform **300** may be a multi-tenant platform instantiated using one or more servers at one or more datacenters including,

for example, the application server **108**, the database server **110**, and the telephony server **112** of the datacenter **106** shown in FIG. 1.

(48) The software platform **300** includes software services accessible using one or more clients. For example, a customer **302** as shown includes four clients—a desk phone **304**, a computer **306**, a mobile device **308**, and a shared device **310**. The desk phone **304** is a desktop unit configured to at least send and receive calls and includes an input device for receiving a telephone number or extension to dial to and an output device for outputting audio and/or video for a call in progress. The computer **306** is a desktop, laptop, or tablet computer including an input device for receiving some form of user input and an output device for outputting information in an audio and/or visual format. The mobile device **308** is a smartphone, wearable device, or other mobile computing aspect including an input device for receiving some form of user input and an output device for outputting information in an audio and/or visual format. The desk phone **304**, the computer **306**, and the mobile device **308** may generally be considered personal devices configured for use by a single user. The shared device **310** is a desk phone, a computer, a mobile device, or a different device which may instead be configured for use by multiple specified or unspecified users.

(49) Each of the clients **304** through **310** includes or runs on a computing device configured to access at least a portion of the software platform **300**. In some implementations, the customer **302** may include additional clients not shown. For example, the customer **302** may include multiple clients of one or more client types (e.g., multiple desk phones or multiple computers) and/or one or more clients of a client type not shown in FIG. 3 (e.g., wearable devices or televisions other than as shared devices). For example, the customer **302** may have tens or hundreds of desk phones, computers, mobile devices, and/or shared devices.

(50) The software services of the software platform **300** generally relate to communications tools, but are in no way limited in scope. As shown, the software services of the software platform **300** include telephony software **312**, conferencing software **314**, messaging software **316**, and other software **318**. Some or all of the software **312** through **318** uses customer configurations **320** specific to the customer **302**. The customer configurations **320** may, for example, be data stored within a database or other data store at a database server, such as the database server **110** shown in FIG. 1.

(51) The telephony software **312** enables telephony traffic between ones of the clients **304** through **310** and other telephony-enabled devices, which may be other ones of the clients **304** through **310**, other VOIP-enabled clients of the customer **302**, non-VOIP-enabled devices of the customer **302**, VOIP-enabled clients of another customer, non-VOIP-enabled devices of another customer, or other VOIP-enabled clients or non-VOIP-enabled devices. Calls sent or received using the telephony software **312** may, for example, be sent or received using the desk phone **304**, a softphone running on the computer **306**, a mobile application running on the mobile device **308**, or using the shared device **310** that includes telephony features.

(52) The telephony software **312** further enables phones that do not include a client application to connect to other software services of the software platform **300**. For example, the telephony software **312** may receive and process calls from phones not associated with the customer **302** to route that telephony traffic to one or more of the conferencing software **314**, the messaging software **316**, or the other software **318**.

(53) The conferencing software **314** enables audio, video, and/or other forms of conferences between multiple participants, such as to facilitate a conference between those participants. In some cases, the participants may all be physically present within a single location, for example, a conference room, in which the conferencing software **314** may facilitate a conference between only those participants and using one or more clients within the conference room. In some cases, one or more participants may be physically present within a single location and one or more other participants may be remote, in which the conferencing software **314** may facilitate a conference between all of those participants using one or more clients within the conference room and one or

more remote clients. In some cases, the participants may all be remote, in which the conferencing software **314** may facilitate a conference between the participants using different clients for the participants. The conferencing software **314** can include functionality for hosting, presenting scheduling, joining, or otherwise participating in a conference. The conferencing software **314** may further include functionality for recording some or all of a conference and/or documenting a transcript for the conference.

(54) The messaging software **316** enables instant messaging, unified messaging, and other types of messaging communications between multiple devices, such as to facilitate a chat or other virtual conversation between users of those devices. The unified messaging functionality of the messaging software **316** may, for example, refer to email messaging which includes a voicemail transcription service delivered in email format.

(55) The other software **318** enables other functionality of the software platform **300**. Examples of the other software **318** include, but are not limited to, device management software, resource provisioning and deployment software, administrative software, third party integration software, and the like. In one particular example, the other software **318** can include a metastore manager, as described herein.

(56) The software **312** through **318** may be implemented using one or more servers, for example, of a datacenter such as the datacenter **106** shown in FIG. **1**. For example, one or more of the software **312** through **318** may be implemented using an application server, a database server, and/or a telephony server, such as the servers **108** through **112** shown in FIG. **1**. In another example, one or more of the software **312** through **318** may be implemented using servers not shown in FIG. **1**, for example, a meeting server, a web server, or another server. In yet another example, one or more of the software **312** through **318** may be implemented using one or more of the servers **108** through **112** and one or more other servers. The software **312** through **318** may be implemented by different servers or by the same server.

(57) Features of the software services of the software platform **300** may be integrated with one another to provide a unified experience for users. For example, the messaging software **316** may include a user interface element configured to initiate a call with another user of the customer **302**. In another example, the telephony software **312** may include functionality for elevating a telephone call to a conference. In yet another example, the conferencing software **314** may include functionality for sending and receiving instant messages between participants and/or other users of the customer **302**. In yet another example, the conferencing software **314** may include functionality for file sharing between participants and/or other users of the customer **302**. In some implementations, some or all of the software **312** through **318** may be combined into a single software application run on clients of the customer, such as one or more of the clients **304** through **310**.

(58) Data used, generated, and/or otherwise interacted with by the software platform may be ingested, stored, and/or processed using an RDBMS. SQL is a programming language used for managing and manipulating relational databases. SQL enables users to interact with databases by defining, querying, and manipulating data stored in tables. SQL is widely used in the field of database management and is supported by most RDBMSs. Some features of SQL include the following. (1) Data Querying: SQL allows users to retrieve specific data from a database using the SELECT statement. Queries can involve filtering, sorting, grouping, and aggregating data. (2) Data Manipulation: SQL provides statements like INSERT, UPDATE, and DELETE for modifying data stored in tables. (3) Data Definition: SQL supports statements like CREATE, ALTER, and DROP for defining and modifying the structure of a database schema, including tables, indexes, views, and constraints. (4) Data Control: SQL includes commands for managing access to the database, such as GRANT and REVOKE, which control privileges like SELECT, INSERT, UPDATE, DELETE, and others.

(59) Before SQL queries are executed by an execution engine, the corresponding SQL code must

be compiled. The terms “compiling” or “compiler” in the context of SQL can encompass all or most of the preprocessing steps that may be necessary prior to execution of SQL code, including the following. (1) Parsing: Parsing involves analyzing the syntactic structure of SQL code to create an abstract representation, typically a parse tree or AST. (2) Validation: Validation ensures that the parsed SQL code conforms to the semantic rules and constraints of the underlying data model, checking elements such as table and column existence and data types. (3) Compilation: Compilation translates the optimized SQL code into a logical plan, which outlines the sequence of operations needed to fulfill the specified operations on the database.

(60) After compilation, an execution engine transforms a logical plan into an execution plan (sometimes called a physical execution plan). The execution plan outlines the detailed steps and operations that the execution engine will perform to execute the computation(s) described by the logical plan. It considers factors such as data partitioning, task scheduling, and resource allocation to efficiently execute the computation(s) across a cluster of computational devices. The execution engine may perform the following key functions. (1) Task Scheduling: The execution engine breaks down the execution plan into smaller tasks and schedules these tasks to run on the available resources within a computer cluster. The execution engine ensures that these tasks are scheduled efficiently across the cluster. (2) Distributed Processing: The execution engine coordinates the distributed processing of data across the nodes, or computers, of the computer cluster. It manages data partitioning, data shuffling, and data movement between nodes to ensure that the computation is performed in parallel and efficiently utilizes the available resources. (3) Fault Tolerance: The execution engine incorporates fault tolerance mechanisms to handle failures gracefully. It monitors the execution of tasks and tracks the lineage of Resilient Distributed Datasets (RDDs) or DataFrames to enable recomputation of lost or corrupted data partitions in case of node failures. (4) Data Caching and Persistence: The execution engine supports caching and persistence of intermediate data to optimize performance. It enables users to specify which data should be cached in memory or persisted to disk, reducing the need for recomputation and improving overall query performance. (4) Integration with Storage Systems: The execution engine interacts with storage systems like Hadoop Distributed File System (HDFS), Amazon Simple Storage Service (S3), or other cloud storage solutions to read input data and write output data generated during query execution. Overall, the execution engine plays a central role in orchestrating the distributed execution of SQL queries and other computations, ensuring efficient resource utilization, fault tolerance, and high-performance data processing on large-scale datasets.

(61) Two common distributed computing systems, or platforms, for compiling and executing SQL code are Apache Spark and Apache Flink. Apache Spark is an open-source distributed computing system that provides an interface for programming entire clusters with implicit data parallelism and fault tolerance. Spark's core abstraction is the RDD, a distributed collection of objects. Spark offers several higher-level libraries built on top of its core, including Spark SQL, which provides support for executing SQL queries (i.e., SQL code) and working with structured data using DataFrame application programming interface (API) or SQL queries. With Spark SQL, users can write SQL queries to process large volumes of data efficiently in parallel across a cluster. Apache Flink is another open-source distributed computing system, which operates according to a data-stream processing framework that provides data distribution, communication, and fault tolerance for distributed computations. Flink's core abstraction is the DataStream API for processing unbounded streams of data and the DataSet API for batch processing. Flink also offers support for SQL queries through Flink SQL, which allows users to write SQL queries to process both streaming and batch data. Flink's SQL support includes a rich set of built-in functions, windowing capabilities, and support for complex event processing (CEP). Although the example implementations described herein focus on Apache Spark, utilization of other distributed computing systems, such as Apache Flink, is within the scope of this disclosure.

(62) In Spark, for example, a logical plan and/or an execution plan, that results from compiled SQL

code, can be represented by DAG at execution time. A DAG is essentially a graph of the computational dependencies between different RDDs or DataFrames resulting from transformations in the SQL code. Each vertex in the DAG represents an operation or transformation, and the edges represent the flow of data between these operations. Both vertices and edges can represent operations in Spark, which are referred to herein as “operational vertices” and “operational edges,” respectively. Traditionally, the DAG produced by Spark is local rather than global. A local DAG represents only a portion of the logical plan and/or execution plan, whereas a global DAG represents an entire logical plan and/or execution plan that results from the compiled SQL code. Spark's DAG is local because Spark constructs the DAG dynamically (e.g., incrementally) just prior to execution of a next stage of the logical plan based on the sequence of transformations and actions invoked by the SQL code. Each RDD or DataFrame transformation adds to the existing DAG, and actions trigger the execution of the DAG up to that point.

(63) FIG. 4 is a block diagram of an example of a Spark execution system **400** for executing SQL queries. The Spark execution system **400** comprises a driver node **402**, a cluster **440**, and a cluster manager **408**. The cluster **440** comprises one or more cluster nodes, which in this example, are illustrated as a worker node **420** and a worker node **430**. The driver node **402** and the cluster manager **408** are communicatively coupled to each other and to the worker node **420** and the worker node **430**, via one or more connections **410** of a network **442**. The network **442** can be or include, for example, the Internet, a local area network (LAN), a wide area network (WAN), a virtual private network (VPN), or another public or private means of electronic computer communication capable of transferring data between computational devices. In some implementations, a computational device can connect to the network **442** via a communal connection point, link, or path, or using a distinct connection point, link, or path. For example, a connection point, link, or path can be wired (e.g., electrical or optical), wireless (e.g., electromagnetic or optical), use other communications technologies, or a combination thereof.

(64) The driver node **402** serves as the central orchestrator of the Spark execution system **400**, comprising a driver program **404** that includes Spark context **406**. The driver program initiates and coordinates the execution of Spark applications, including the processing of SQL queries. The driver program **404** interfaces with the Spark context **406** to establish communication with the worker nodes **420** and **430**, to initialize the execution environment, and to oversee task distribution and resource allocation. The cluster manager **408** is responsible for managing and coordinating the worker nodes **420** and **430** within the cluster **440**.

(65) Each worker node **420** and **430** comprises an executor **422** and **432**, respectfully, which is responsible for executing tasks delegated by the driver node **402**. In some implementations, each executor **422** and **432** is a processor, such as a microprocessor. The executors **422** and **432** implement respective caches **426** and **436**, for efficient data storage and retrieval during computation. Furthermore, each executor **422** and **432** implement one or more respective cores **424** and **434**, which are responsible for carrying out specific computations as part of an SQL query execution process.

(66) FIG. 5 is an example of a Spark optimization framework **500** for translating one or more SQL queries **506** of SQL code into executable operations on a database. The framework **500** can be partitioned into a frontend **502**, a compiler **510**, and a backend **518**. The frontend **502** serves as the entry point for users to provide the SQL query **506**, as well as other instruction formats such as a Hive query **504** and a DataFrame **508**. The frontend **502** is responsible for accepting the various input formats and passing them to the compiler **510** for further processing.

(67) The compiler **510** is the primary component responsible for translating SQL query **506** (and/or other inputs) into an optimized logical plan **516**. The compiler **510** processes the SQL query **506** in several stages. The first stage yields an unresolved logical plan **512**, which represents the initial representation of the SQL query **506** without any optimization applied. The second stage yields a logical plan **514**, which is a refinement of the unresolved logical plan **512** based on resolving

dependencies and applying standard optimization techniques. The third stage yields an optimized logical plan **516**, which is a further refinement of the logical plan **514** based on advanced optimization strategies to further enhance query performance and efficiency. The unresolved logical plan **512**, the logical plan **514**, and/or the optimized logical plan **516** may each be represented by a local DAG, i.e., a DAG representing the SQL query **506** which is a part of more comprehensive SQL code or an SQL job.

(68) Once the compiler **510** generates the optimized logical plan **516**, it passes the optimized logical plan **516** to the backend **518** for execution. The backend **518** comprises several stages that yield, respectively, one or more initial physical plans **520**, a selected physical plan **524**, and RDDs **546**. The initial physical plans **520** represent the actual execution plans generated from the optimized logical plan **516** according to various resource allocations. The cost model **522** evaluates the cost associated with the different initial physical plans **520**, considering factors such as data distribution, resource utilization, and execution time. Based on the analysis of the cost model **522**, the backend **518** selects the most efficient of the initial physical plans **520**, referred to as the selected physical plan **524**. Finally, the backend **518** translates the selected physical plan **524** into RDDs for distributed data processing, which enables parallel execution of SQL queries across multiple nodes in a cluster, such as the worker node **420** and the worker node **430** in the cluster **440** of FIG. **4**. The distributed execution of SQL queries enhances scalability and performance, particularly for handling large volumes of data.

(69) FIG. **6** is a diagram illustrating an example of a data management environment **600**. The data management environment **600** includes a first data management platform **602**, a second data management platform **604**, and a data processing pipeline **606**. The first data management platform **602** includes a first data processing platform **608**, a first metastore **610** associated with the first data processing platform **608**, and a first data store associated with the first data processing platform **608**. Similarly, the second data management platform **604** includes a second data processing platform **614**, a second metastore **616** associated with the second data processing platform **614**, and a second data store **618** associated with the second data processing platform **614**. The first metastore **610** maintains metadata associated with data stored in the first data store **612**. The second metastore **616** maintains metadata associated with data stored in the second data store **618**. A metastore manager **620** associated with the second data management platform **604** may be configured to manage interactions between the first data management platform **602** and the second data management platform **604**. The metastore manager **620** also may be configured to manage interactions between the second data management platform **604** and the data processing pipeline **606**. A refresh component **622** is maintained on the data processing pipeline **606** to facilitate access, by the data processing pipeline **606**, to data stored in the first data store **612**.

(70) FIG. **7** is a block diagram of an example of a metastore manager system. A metastore manager **700**, which in some cases may be implemented in connection with a software platform (e.g., the software platform **300** shown in FIG. **3**), is accessed by a driver node **702** (e.g., the driver node **402** shown in FIG. **4**) and used to facilitate interactions between data management platforms. The metastore manager **700** may be implemented using one or more servers and software running thereon. For example, the metastore manager **700** may be implemented using one or more of the servers **108** through **112** shown in FIG. **1** and may use communication software such as or similar to the software **312** through **318** shown in FIG. **3**. In some examples, the metastore manager **700** may be implemented using the driver node **702** and/or one or more worker nodes (e.g., the worker node **420** and/or the worker node **430** shown in FIG. **4**). The metastore manager **700** may include, or be included in, a conferencing system and/or the contact center. The metastore manager **700** may be, be similar to, include, or be included in, the computing device **200** shown in FIG. **2**. The metastore manager **700** includes software for facilitating exchange of data between two or more of a data management platform **704**, a data management platform **706**, and a data processing pipeline **708**. As shown, the software includes access processing software **710**, synchronization processing

software **712**, and data exchange software **714**.

(71) The access processing software **710** facilitates access to a data store and/or a metastore to obtain data for a data processing operation. For example, the driver node **702** may forward a request and/or a function call to the access processing software **710** to facilitate obtaining data from a data management platform **704** or **706**. The driver node **702** may be associated one or more of the data management platforms **704** or **706**. The access processing software **710** may determine and otherwise manage permissions, roles, and/or other data access attributes, initiate a synchronization operation (e.g., by causing an activation of the synchronization processing software **712**), and/or provide access to the requested and/or called data.

(72) The synchronization processing software **712** may facilitate synchronization operations between two or more of the data management platforms **704** or **706**. For example, in some cases, the synchronization processing software **712** may be configured to facilitate periodic synchronization operations and/or on-demand synchronization operations. A periodic synchronization operation may be a synchronization operation that is facilitated according to a schedule that may be maintained by the metastore manager **700**. An on-demand synchronization operation may be a synchronization operation that is facilitated in response to a request or function call received by the access processing software **710**.

(73) The data exchange software **714** may facilitate the transfer of data from one data management platform to another data management platform. The data exchange software **714** may include any number of different security components, query engines, data transformation components, and/or other components configured to facilitate providing data stored in one data management platform to another data management platform.

(74) Although the access processing software **710**, synchronization processing software **712**, and data exchange software **714** are shown as separate software components, in some implementations, some or all of the access processing software **710**, synchronization processing software **712**, and data exchange software **714** may be combined. For example, the metastore manager **700** may be or include a single software component which performs the functionality of all of the access processing software **710**, synchronization processing software **712**, and data exchange software **714**. In some implementations, one or more of the access processing software **710**, synchronization processing software **712**, and data exchange software **714** may be comprised of multiple software components. In some implementations, the metastore manager **700** may include software components other than the access processing software **710**, synchronization processing software **712**, and data exchange software **714**, such as in addition to or in place of one or more of the access processing software **710**, synchronization processing software **712**, and data exchange software **714**.

(75) FIG. **8** is a schematic block diagram illustrating an example **800** of an implementation of a multi-platform data management environment in accordance with the present disclosure. The example **800** includes a first data processing platform **802**, a metastore manager **804**, and a data processing pipeline **806**. As shown, the first data processing platform **802** may perform a data processing operation **808** in association with a set **810** of data maintained in a first data store **812** associated with the first data processing platform **802**. The first data processing platform **802** and the first data store **812** may be associated with a first data management platform (e.g., the first data management platform **602** depicted in FIG. **6**). The first data processing platform **802** may obtain the set **810** of data from the first data store **812** via a first metastore **814**. The first data processing platform **802** may generate, by performing the data processing operation **808**, a set **816** of processed data (shown as “x” in FIG. **8**).

(76) The first data processing platform **802** may store the set **816** of processed data in the first data store **812**. The first data processing platform **802** may store a first set **818** of metadata corresponding to the set **816** of processed data in the first metastore **814**. The first set **818** of metadata may include partition metadata **820** indicative of partitioning information associated with

the set **816** of processed data. The metastore manager **804** may store, in a second metastore **822** associated with the second data management platform, a second set **824** of metadata corresponding to a subset of the first set **818** of metadata. The second set **824** of metadata may exclude the partition metadata **820**. The metastore manager **804** may cause the data processing pipeline **806** to execute a refresh component **826** to obtain the partition metadata **820**. For example, the data processing pipeline **806** may execute the refresh component **826**, causing the first metastore **814** (e.g., via the first data processing platform **802**) to provide the partition metadata **820** to the data processing pipeline **806**. The metastore manager **804** may obtain, in association with the partition metadata **820**, the set **816** of processed data using the data processing pipeline **806**.

(77) The data processing pipeline **806** may be communicatively coupled, via a platform interface, with the second data management platform. The second data management platform may include a second data store **828** associated with the second metastore **822**. The data processing pipeline **806** may not, in some implementations, be communicatively coupled with the first data management platform. The data processing operation **808** may include a SQL query on a distributed database associated with at least one of the first data store **812** or the second data store **828**. The metastore manager **804** may obtain a directed acyclic graph (DAG) corresponding to the data processing operation **808** and may configure the data processing operation **808** in association with the DAG. The metastore manager **804** may configure the data processing operation **808** at least in part by allocating a set of computational resources for the data processing operation **808**. The data processing pipeline **806** may perform an additional data processing operation in association with the set **816** of processed data.

(78) The data processing pipeline **806** may include an artificial intelligence (AI) pipeline such as, for example, a machine learning pipeline. The AI pipeline may be used to generate content, predict events, make decisions, and/or otherwise process information. In some examples, the AI pipeline may take customer data, communication data, and/or environmental data as input. Customer data may include any number of different types of data associated with customers of the entity that is using the data processing pipeline **806**. Data associated with a customer may include, for example, an account number, a name, an address, an email address, demographic information, a phone number, an internet protocol (IP) address, a customer preference, an attribute of a service and/or a product that the customer has paid for, among other examples. Communication data may include any number of different types of data associated with communications that have occurred (e.g., involving customers). Data associated with a communication may include, for example, device identifiers, domain names, quality of service (QoS) requirements, QoS rules, identifications of communications types (e.g., audio, video, workspace sharing), identifications of communication quality, identifications of communication media (e.g., wired networks, wireless networks), identifications of communication protocols, attributes associated with communication devices, network traffic information, and/or communication durations, among other examples. Environmental data may include any number of different types of contextual data associated with customers and/or communications. Contextual data associated with customers and/or communications may include market information, societal information, local information associated with a communication device (e.g., activation times, use durations, temperatures, device capabilities, and/or software versions among other examples).

(79) In some cases, rather than including a refresh component in the data processing pipeline to facilitate obtaining data generated by a first data processing operation, a data processing operation that is configured in association with the first data management platform may be submitted to the second data management platform to be performed, so that the resulting output is stored by the second data management platform.

(80) FIG. **9** is a schematic block diagram illustrating another example **900** of an implementation of a multi-platform data management environment in accordance with the present disclosure. The example **900** includes a first data processing platform **902**, a second data processing platform **904**,

and a data processing pipeline **906**. The first data processing platform **902** may be associated with a first data management platform (e.g., the data management platform **704** shown in FIG. 7 or the first data management platform **602** shown in FIG. 6) and the second data processing platform **904** may be associated with a second data management platform (e.g., the data management platform **706** or **708** shown in FIG. 7 or the second data management platform **604** shown in FIG. 6). As shown, the first data processing platform **902** may provide a job submission **908** to the second data processing platform **904**. In some examples, the job submission **908** may be provided directly to the second data processing platform **904**. In some other examples, the job submission **908** may be provided to the second data processing platform **904** via a metastore manager **910**. For example, the first data processing platform **902** may cause the metastore manager **910** to provide the job submission **908** to the second data processing platform **904**. The job submission **908** may be or include a script that can be executed to perform a data processing operation **912**, a function call to initiate the data processing operation **912**, an entity that describes the data processing operation **912**, and/or a command configured to cause execution of the data processing operation **912**. In some examples, the metastore manager **910** may provide a submission command **914** to the first data processing platform **902** to cause the first data processing platform **902** to provide the job submission **908** to the second data processing platform **904**.

(81) The second data processing platform **904** may perform the data processing operation **912** in association with a set **916** of data maintained in a first data store **918** associated with the first data processing platform **902**. The first data store **918** may be associated with the first data management platform. The second data processing platform **904** may obtain the set **916** of data from the first data store **918** via a first metastore **920** associated with the first data store **918**. The second data processing platform **904** may obtain the set **916** of data from the first data store **918** via a second metastore **922**. The second metastore **922** may be associated with the second data processing platform **904** and a second data store **924**. The second metastore **922** and the second data store **924** may be associated with the second data management platform. The second data processing platform **904** may generate, by performing the data processing operation **912**, a set **926** of processed data (shown as “x” in FIG. 9). The data processing operation **912** may include a SQL query on a distributed database associated with the first data store **918**. The second data processing platform **904** may store the set **926** of processed data in the second data store **924**. The second data processing platform **904** may store a set **928** of metadata corresponding to the set **926** of processed data in the second metastore **922**. The set **928** of metadata may include partition metadata **930** indicative of partitioning information associated with the set **926** of processed data. The data processing pipeline **906** may be communicatively coupled, via a platform interface, with the second data management platform. For example, the data processing pipeline **906** may be communicatively coupled with the second metastore **922**. The data processing pipeline **906** may perform an additional data processing operation in association with the set **926** of processed data.

(82) To further describe some implementations in greater detail, reference is next made to examples of techniques which may be performed by or using a system for performing multi-platform data operations. FIG. 10 is a flowchart of an example of a technique **1000** for performing a multi-platform data operation. The technique **1000** can be executed using computing devices, such as the systems, hardware, and software described with respect to FIGS. 1-9. The technique **1000** can be performed, for example, by executing a machine-readable program or other computer-executable instructions, such as routines, instructions, programs, or other code. The steps, or operations, of the technique **1000**, or another technique, method, process, or algorithm described in connection with the implementations disclosed herein can be implemented directly in hardware, firmware, software executed by hardware, circuitry, or a combination thereof.

(83) For simplicity of explanation, the technique **1000** is depicted and described herein as a series of steps or operations. However, the steps or operations of the technique **1000** in accordance with this disclosure can occur in various orders and/or concurrently. Additionally, other steps or

operations not presented and described herein may be used. Furthermore, not all illustrated steps or operations may be required to implement a technique in accordance with the disclosed subject matter.

(84) The step **1002** includes initiating, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation to generate a first processed data set. The first processed data set is stored in the first data store and a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store. The first metadata set includes partition metadata indicative of partitioning information associated with the first processed data set. The step **1004** includes initiating a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set. The step **1006** includes activating a data processing pipeline. The data processing pipeline may be configured to perform one or more of a number of steps such as, for example, one or more of steps **1008**, **1010**, **1012**, and **1014**.

(85) The step **1008** includes obtaining the partition metadata using a refresh component. The step **1010** includes obtaining, based on the partition metadata, the first processed data set. The step **1012** includes performing a second data processing operation on the first processed data set to generate a second processed data set. The step **1014** includes outputting the second processed data set. The data processing pipeline may be communicatively coupled, via a platform interface, with the second data management platform. In some examples, the data processing pipeline may be communicatively coupled, via a platform interface, with only the second data management platform. The first data processing operation may include an SQL query on a distributed database associated with at least one of the first data store or a second data store associated with the second metastore. In some examples, the method may include obtaining a DAG corresponding to the first data processing operation and configuring the first data processing operation in association with the DAG. For example, configuring the first data processing operation may include configuring the first data processing operation at least in part by allocating a set of computational resources for the first data processing operation. The second processed data set may include an output of an AI operation. The data processing pipeline, to output the second processed data set, may be configured to output the second processed data set to a UCaaS platform.

(86) Some implementations include a method, comprising: initiating, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation to generate a first processed data set, wherein the first processed data set is stored in the first data store, and wherein a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store, the first metadata set comprising partition metadata indicative of partitioning information associated with the first processed data set; initiating a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set; and activating a data processing pipeline configured to: obtain the partition metadata using a refresh component; obtain, based on the partition metadata, the first processed data set; perform a second data processing operation on the first processed data set to generate a second processed data set; and output the second processed data set.

(87) In some implementations, the data processing pipeline is communicatively coupled, via a platform interface, with the second data management platform.

(88) In some implementations, the data processing pipeline is communicatively coupled, via a platform interface, with only the second data management platform.

(89) In some implementations, the first data processing operation comprises an SQL query on a distributed database associated with at least one of the first data store or a second data store associated with the second metastore.

(90) In some implementations, the method further comprises: obtaining a DAG corresponding to

the first data processing operation; and configuring the first data processing operation in association with the DAG.

(91) In some implementations, the method further comprises: configuring the first data processing operation at least in part by allocating a set of computational resources for the first data processing operation.

(92) In some implementations, the second processed data set comprises an output of an AI operation.

(93) In some implementations, to output the second processed data set, the data processing pipeline is configured to output the second processed data set to a UCaaS platform.

(94) Some implementations include a non-transitory computer-readable medium storing instructions operable to cause one or more processors to perform operations comprising: initiating, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation to generate a first processed data set, wherein the first processed data set is stored in the first data store, and wherein a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store, the first metadata set comprising partition metadata indicative of partitioning information associated with the first processed data set; initiating a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set; and activating a data processing pipeline configured to: obtain the partition metadata using a refresh component; obtain, based on the partition metadata, the first processed data set; perform a second data processing operation on the first processed data set to generate a second processed data set; and output the second processed data set.

(95) In some implementations, the data processing pipeline is configured to obtain the first processed data set via a platform interface that connects the data processing pipeline with only the second data management platform.

(96) In some implementations, the first data processing operation comprises an SQL query on a distributed database associated with at least one of the first data store or a second data store associated with the second metastore.

(97) In some implementations, the operations further comprising allocating a set of computational resources for the first data processing operation.

(98) In some implementations, the data processing pipeline comprises an AI pipeline.

(99) Some implementations include a system, comprising: one or more memories; and one or more processors configured to execute instructions stored in the one or more memories to: initiate, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation to generate a first processed data set, wherein the first processed data set is stored in the first data store, and wherein a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store, the first metadata set comprising partition metadata indicative of partitioning information associated with the first processed data set; initiate a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set; and activate a data processing pipeline configured to: obtain the partition metadata using a refresh component; obtain, based on the partition metadata, the first processed data set; perform a second data processing operation on the first processed data set to generate a second processed data set; and output the second processed data set.

(100) In some implementations, the system further comprises a platform interface that communicatively couples the data processing pipeline with the second data management platform.

(101) In some implementations, the system further comprises a platform interface that communicatively couples the data processing pipeline with only the second data management platform.

(102) In some implementations, the system further comprises a platform interface that communicatively couples the first data management platform with the second data management platform.

(103) In some implementations, the system further comprises a platform interface that communicatively couples the data processing pipeline with a UCaaS platform.

(104) In some implementations, the second data management platform comprises a metastore manager configured to manage access, by the data processing pipeline, to the first data store.

(105) In some implementations, the second data processing operation comprises an AI operation.

(106) As used herein, unless explicitly stated otherwise, any term specified in the singular may include its plural version. For example, “a computer that stores data and runs software,” may include a single computer that stores data and runs software or two computers—a first computer that stores data and a second computer that runs software. Also “a computer that stores data and runs software,” may include multiple computers that together stored data and run software. At least one of the multiple computers stores data, and at least one of the multiple computers runs software.

(107) As used herein, the term “computer-readable medium” encompasses one or more computer readable media. A computer-readable medium may include any storage unit (or multiple storage units) that store data or instructions that are readable by processing circuitry. A computer-readable medium may include, for example, at least one of a data repository, a data storage unit, a computer memory, a hard drive, a disk, or a random access memory. A computer-readable medium may include a single computer-readable medium or multiple computer-readable media. A computer-readable medium may be a transitory computer-readable medium or a non-transitory computer-readable medium.

(108) As used herein, the term “memory subsystem” includes one or more memories, where each memory may be a computer-readable medium. A memory subsystem may encompass memory hardware units (e.g., a hard drive or a disk) that store data or instructions in software form. Alternatively or in addition, the memory subsystem may include data or instructions that are hard-wired into processing circuitry.

(109) As used herein, processing circuitry includes one or more processors. The one or more processors may be arranged in one or more processing units, for example, a central processing unit (CPU), a graphics processing unit (GPU), or a combination of at least one of a CPU or a GPU.

(110) As used herein, the term “engine” may include software, hardware, or a combination of software and hardware. An engine may be implemented using software stored in the memory subsystem. Alternatively, an engine may be hard-wired into processing circuitry. In some cases, an engine includes a combination of software stored in the memory subsystem and hardware that is hard-wired into the processing circuitry.

(111) The implementations of this disclosure can be described in terms of functional block components and various processing operations. Such functional block components can be realized by a number of hardware or software components that perform the specified functions. For example, the disclosed implementations can employ various integrated circuit components (e.g., memory elements, processing elements, logic elements, look-up tables, and the like), which can carry out a variety of functions under the control of one or more microprocessors or other control devices. Similarly, where the elements of the disclosed implementations are implemented using software programming or software elements, the systems and techniques can be implemented with a programming or scripting language, such as C, C++, Java, JavaScript, assembler, or the like, with the various algorithms being implemented with a combination of data structures, objects, processes, routines, or other programming elements.

(112) Functional aspects can be implemented in algorithms that execute on one or more processors. Furthermore, the implementations of the systems and techniques disclosed herein could employ a number of conventional techniques for electronics configuration, signal processing or control, data processing, and the like. The words “mechanism” and “component” are used broadly and are not

limited to mechanical or physical implementations, but can include software routines in conjunction with processors, etc. Likewise, the terms “system” or “tool” as used herein and in the figures, but in any event based on their context, may be understood as corresponding to a functional unit implemented using software, hardware (e.g., an integrated circuit, such as an ASIC), or a combination of software and hardware. In certain contexts, such systems or mechanisms may be understood to be a processor-implemented software system or processor-implemented software mechanism that is part of or callable by an executable program, which may itself be wholly or partly composed of such linked systems or mechanisms.

(113) Implementations or portions of implementations of the above disclosure can take the form of a computer program product accessible from, for example, a computer-usable or computer-readable medium. A computer-usable or computer-readable medium can be a device that can, for example, tangibly contain, store, communicate, or transport a program or data structure for use by or in connection with a processor. The medium can be, for example, an electronic, magnetic, optical, electromagnetic, or semiconductor device.

(114) Other suitable mediums are also available. Such computer-usable or computer-readable media can be referred to as non-transitory memory or media, and can include volatile memory or non-volatile memory that can change over time. The quality of memory or media being non-transitory refers to such memory or media storing data for some period of time or otherwise based on device power or a device power cycle. A memory of an apparatus described herein, unless otherwise specified, does not have to be physically contained by the apparatus, but is one that can be accessed remotely by the apparatus, and does not have to be contiguous with other memory that might be physically contained by the apparatus.

(115) While the disclosure has been described in connection with certain implementations, it is to be understood that the disclosure is not to be limited to the disclosed implementations but, on the contrary, is intended to cover various modifications and equivalent arrangements included within the scope of the appended claims, which scope is to be accorded the broadest interpretation so as to encompass all such modifications and equivalent structures as is permitted under the law.

Claims

1. A method, comprising: communicating a request, from a driver node to a metastore manager, for initiating, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation, to be performed by the first data management platform, to generate a first processed data set, wherein the first processed data set is stored in the first data store, and wherein a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store, the first metadata set comprising partition metadata indicative of partitioning information associated with the first processed data set, wherein the metastore manager is configured to: obtain a directed acyclic graph (DAG) corresponding to the first data processing operation; and configure the first data processing operation in association with the DAG, wherein the DAG includes an unresolved logical plan, a logical plan, and an optimized logical plan; initiating, by the metastore manager, a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set; and activating, by the metastore manager, a data processing pipeline that is not communicatively coupled with the first data management platform wherein the data processing pipeline includes an artificial intelligence (AI) pipeline, to cause the data processing pipeline to obtain the partition metadata using a refresh component, and wherein the data processing pipeline is configured to: obtain, based on the partition metadata, the first processed data set; perform a second data processing operation on the first processed data set to generate a second processed data set, wherein the second processed data set comprises an output of an AI operation; and output the second processed data set to a unified

communications as a service (UCaaS) platform.

2. The method of claim 1, wherein the data processing pipeline is communicatively coupled, via a platform interface, with the second data management platform.

3. The method of claim 1, wherein the data processing pipeline is communicatively coupled, via a platform interface, with only the second data management platform.

4. The method of claim 1, wherein the first data processing operation comprises a structured query language (SQL) query on a distributed database associated with at least one of the first data store or a second data store associated with the second metastore.

5. The method of claim 1, further comprising configuring the first data processing operation at least in part by allocating a set of computational resources for the first data processing operation.

6. A non-transitory computer-readable medium storing instructions operable to cause one or more processors to perform operations comprising: communicating a request, from a driver node to a metastore manager, for initiating, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation, to be performed by the first data management platform, to generate a first processed data set, wherein the first processed data set is stored in the first data store, and wherein a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store, the first metadata set comprising partition metadata indicative of partitioning information associated with the first processed data set, wherein the metastore manager is configured to: obtain a directed acyclic graph (DAG) corresponding to the first data processing operation; and configure the first data processing operation in association with the DAG, wherein the DAG includes an unresolved logical plan, a logical plan, and an optimized logical plan; initiating, by the metastore manager, a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set; and activating, by the metastore manager, a data processing pipeline that is not communicatively coupled with the first data management platform wherein the data processing pipeline includes an artificial intelligence (AI) pipeline, to cause the data processing pipeline to: obtain the partition metadata using a refresh component, and wherein the data processing pipeline is configured to: obtain, based on the partition metadata, the first processed data set; perform a second data processing operation on the first processed data set to generate a second processed data set, wherein the second processed data set comprises an output of an AI operation; and output the second processed data set to a unified communications as a service (UCaaS) platform.

7. The non-transitory computer-readable medium of claim 6, wherein the data processing pipeline is configured to obtain the first processed data set via a platform interface that connects the data processing pipeline with only the second data management platform.

8. The non-transitory computer-readable medium of claim 6, wherein the first data processing operation comprises a structured query language (SQL) query on a distributed database associated with at least one of the first data store or a second data store associated with the second metastore.

9. The non-transitory computer-readable medium of claim 6, the operations further comprising allocating a set of computational resources for the first data processing operation.

10. A system, comprising: one or more memories; and one or more processors configured to execute instructions stored in the one or more memories to: communicating a request, from a driver node to a metastore manager, for initiating, in association with a data set maintained in a first data store associated with a first data management platform, a first data processing operation, to be performed by the first data management platform, to generate a first processed data set, wherein the first processed data set is stored in the first data store, and wherein a first metadata set corresponding to the first processed data set is stored in a first metastore associated with the first data store, the first metadata set comprising partition metadata indicative of partitioning information associated with the first processed data set, wherein the metastore manager is configured to: obtain a directed acyclic graph (DAG) corresponding to the first data processing

operation; and configure the first data processing operation in association with the DAG, wherein the DAG includes an unresolved logical plan, a logical plan, and an optimized logical plan; initiate, by the metastore manager, a synchronization operation to store, in a second metastore associated with a second data management platform, a second metadata set corresponding to a subset of the first metadata set; and activate, by the metastore manager, a data processing pipeline that is not communicatively coupled with the first data management platform wherein the data processing pipeline includes an artificial intelligence (AI) pipeline, to cause the data processing pipeline to obtain the partition metadata using a refresh component, and wherein the data processing pipeline is configured to: obtain, based on the partition metadata, the first processed data set; perform a second data processing operation on the first processed data set to generate a second processed data set, wherein the second processed data set comprises an output of an AI operation; and output the second processed data set to a unified communications as a service (UCaaS) platform.

11. The system of claim 10, further comprising a platform interface that communicatively couples the data processing pipeline with the second data management platform.

12. The system of claim 10, further comprising a platform interface that communicatively couples the data processing pipeline with only the second data management platform.

13. The system of claim 10, further comprising a platform interface that communicatively couples the first data management platform with the second data management platform.

14. The system of claim 10, further comprising a platform interface that communicatively couples the data processing pipeline with a unified communications as a service (UCaaS) platform.

15. The system of claim 10, wherein the second data management platform comprises a metastore manager configured to manage access, by the data processing pipeline, to the first data store.
