

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12393641
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Thouppuarachchi; Chirath Neranjena

---

### Methods for utilizing solver hardware for solving partial differential equations

---

#### Abstract

Embodiments relate to a computing system for solving differential equations. The system is configured to receive problem packages corresponding to problems to be solved, each comprising at least a differential equation and a domain, and to select a solver of a plurality of solvers, based upon availability of each of the plurality of solvers. A dispatch computer selects a solver by monitoring the plurality of solvers, and responsive to a solver becoming available, determines if a received problem package having at least a threshold priority level can be solved by the solver. Otherwise, the dispatch computer generates a plurality of solver scenarios each reflecting a permutation of received problem packages assigned to solvers estimated to become available within a threshold period of time, and assigns the problem packages in accordance with a solver scenario having a highest utilization score.

---

<b>Inventors:</b>	<b>Thouppuarachchi; Chirath Neranjena (Redwood City, CA)</b>
<b>Applicant:</b>	<b>Vorticity Inc. (Redwood City, CA)</b>
<b>Family ID:</b>	<b>1000008762366</b>
<b>Assignee:</b>	<b>Vorticity Inc. (Redwood City, CA)</b>
<b>Appl. No.:</b>	<b>18/411841</b>
<b>Filed:</b>	<b>January 12, 2024</b>

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20240211531 A1	Jun. 27, 2024

#### Related U.S. Application Data

## Publication Classification

**Int. Cl.:** **G06F17/13** (20060101); **G06F11/34** (20060101); **G06F30/20** (20200101); G06F111/10 (20200101)

**U.S. Cl.:**

**CPC** **G06F17/13** (20130101); **G06F11/3409** (20130101); **G06F30/20** (20200101); G06F2111/10 (20200101)

## Field of Classification Search

**CPC:** G06F (17/13); G06F (2111/10); G06F (30/20); G06F (11/3409); G06F (30/28)

**USPC:** 708/446

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
4720780	12/1987	Dolecek	N/A	N/A
6041398	12/1999	Pechanek et al.	N/A	N/A
8583720	12/2012	Kusmanoff et al.	N/A	N/A
10151856	12/2017	Copeland	N/A	N/A
2013/0116988	12/2012	Zhang et al.	N/A	N/A
2014/0324399	12/2013	Barron et al.	N/A	N/A
2016/0267349	12/2015	Shoaib et al.	N/A	N/A
2018/0307980	12/2017	Barik et al.	N/A	N/A
2018/0314780	12/2017	Bertilsson et al.	N/A	N/A
2020/0301876	12/2019	Hamlin	N/A	N/A
2021/0019591	12/2020	Venkatesh et al.	N/A	N/A
2022/0148130	12/2021	Tang et al.	N/A	N/A

### OTHER PUBLICATIONS

Evans, D., et al. "Construction of extrapolation tables by systolic arrays for solving ordinary differential equations." Parallel computing vol. 4, Issue 1, Feb. 1, 1987, pp. 33-48. cited by applicant

Evans, D., et al. "Systolic arrays for group explicit methods for solving parabolic partial differential equations." The Computer Society, May 25-27, 1988, pp. 315-329. cited by applicant

Jainandunsing, K. "Parallel algorithms for solving systems of linear equations and their mapping on systolic arrays," College of Deans, Jan. 3, 1989, pp. 1-195. cited by applicant

Strbac, P., et al. "Hierarchical Model of a Systolic Array for Solving Differential Equations Implemented as an Upgraded Petri Net", WSEAS Transactions on Systems, Issue 1, vol. 8, Jan. 2009, pp. 1-10. cited by applicant

United States Office Action, U.S. Appl. No. 16/989,821, filed Jul. 29, 2022, 22 pages. cited by applicant

United States Office Action, U.S. Appl. No. 16/989,829, filed Feb. 11, 2022, 8 pages. cited by applicant  
United States Office Action, U.S. Appl. No. 16/989,829, filed Aug. 30, 2021, 22 pages. cited by applicant  
United States Office Action, U.S. Appl. No. 16/989,825, filed Jun. 28, 2023, 7 pages. cited by applicant  
United States Office Action, U.S. Appl. No. 18/194,321, filed Feb. 15, 2024, 22 pages. cited by applicant

---

*Primary Examiner:* Mai; Tan V

*Attorney, Agent or Firm:* Fenwick & West LLP

---

## **Background/Summary**

CROSS-REFERENCE TO RELATED APPLICATION (1) This application is a continuation of pending U.S. application Ser. No. 16/989,825, filed on Aug. 10, 2020, which claims the benefit of U.S. Provisional Application No. 62/889,550, filed on Aug. 20, 2019, the forgoing which are hereby incorporated by reference in their entireties.

### **BACKGROUND**

- (1) The present disclosure generally relates to methods of using dedicated solver hardware for solving partial differential equations.
- (2) Differential equations are ubiquitous in describing fundamental laws of nature, human interactions and many other phenomena. Applications include fluid dynamics, molecular dynamics, electronic structure, high frequency options trading, brain tissue simulations, satellite orbitals, nuclear explosion simulations, black hole simulations, etc.
- (3) Solving differential equations has been a major use of computers since their advent in the mid-1900s. Today, estimates show that over 50% of high performance computing is diverted towards solving differential equations, from supercomputers at national labs to small computer clusters in medium size companies. As such, a need exists for computers that can more efficiently solve differential equations.

### **SUMMARY**

- (4) Embodiments relate to a system for solving differential equations. The system is configured to receive and enqueue problem packages corresponding to problems to be solved, each comprising at least a differential equation and a domain. A dispatch computer retrieves problem packages from the queue, and selects a solver of a plurality of solvers, based upon availability of each of the plurality of solvers, for which to assign the problem package. In some embodiments, the dispatch computer selects a solver by, responsive to a solver becoming available, determining if a received problem package in the queue has at least a threshold priority level, and if the received problem package can be solved by the solver. If so, the dispatch computer assigns the identified high-priority problem package to the solver.
- (5) In one or more embodiments, if the solver is not available (e.g., if no problem package has at least the threshold priority level or if the available solver is not suitable for solving the problem package), the dispatch computer determines estimated completion times for solvers of the plurality of solvers estimated to become available within a threshold period of time, and generates a plurality of solver scenarios each reflecting a permutation of received problem packages assigned to solvers estimated to become available within the threshold period of time. The dispatch computer

determines a utilization metric for each generated scenario, assigns the problem packages in accordance with a selected solver scenario having a highest utilization metric.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a graph illustrating the effects of Amdahl's law.
- (2) FIG. 2 illustrates a high level diagram of a system for solving differential equations, in accordance with some embodiments.
- (3) FIG. 3 illustrates a diagram of the interface computer, dispatch computer, and solver units, in accordance with some embodiments.
- (4) FIG. 4 is a diagram illustrating components of a solver, in accordance with some embodiments.
- (5) FIG. 5 is a flowchart of a process for solving a differential equation, in accordance with some embodiments.
- (6) FIG. 6 is a flowchart illustrating a process for dispatching problems to solvers, in accordance with some embodiments.
- (7) FIG. 7 illustrates a layout of a Differential Equation Accelerator (DEA), in accordance with some embodiments.
- (8) FIG. 8 is a flowchart of a process for using a DEA to solve a subdomain, in accordance with some embodiments.
- (9) The figures depict embodiments of the present disclosure for purposes of illustration only. One skilled in the art will readily recognize from the following description that alternative embodiments of the structures and methods illustrated herein may be employed without departing from the principles, or benefits touted, of the disclosure described herein.

### DETAILED DESCRIPTION

- (10) Embodiments herein are directed to a purpose built computing architecture to enable fast solving of differential equations within large domains with complicated boundary conditions. Differential equations are ubiquitous in describing fundamental laws of nature, human interactions and many other phenomena. Applications of differential equations include fluid dynamics, molecular dynamics, electronic structure, high frequency options trading, brain tissue simulations, satellite orbitals, nuclear explosion simulations, black hole simulations, etc.
- (11) While simple differential equations can be solved with analytical solutions, many more complicated differential equations must be solved numerically in order to obtain useful results. This usually involves breaking up a problem domain into many slices/nodes/particles etc., and solving a discretized form of the equation on each slice/node/particle. This can be a tedious process. In addition, as domain size and accuracy requirements increase (e.g., resolution of the solution, maximum partition size possible, etc.), the number of calculations needed to be performed can increase dramatically.
- (12) The usage of current computer systems (e.g., general-purpose computers) has several problems. In many applications, each particle or node in the domain of the differential equation to be solved requires perhaps  $\sim 10^{2-3}$  floating point operations to calculate the next time step. Since these operations have to be done sequentially, the best time scaling that the simulation or solution can achieve is described in Equation (1) below, even without accounting for the clock cycles needed in a von Neumann architecture to fetch instructions, decode, access memory multiple times to perform a single operation.

$$(13) \quad \frac{\text{computer\_clockspeed}[\text{s}^{-1}]}{\text{operations\_per\_node}[\text{timestep}^{-1}] \times \text{number\_of\_nodes\_in\_domain}} \text{timesteps / second} \quad (1)$$

- (14) For problems that require “strong vertical scaling” such as molecular dynamics, this ceiling is a major problem where even the best supercomputers can only muster several microseconds of

simulation time for several days' worth of compute time.

(15) When possible, for large domain sizes, parallel computing can be used to speed up the solution. However the need to pass large amounts of data between the computing units in such a setup slows down the time to solution. For example, when engineers use 1000 cores, the speed up is no more than 10 times as using a single core. This problem is generally referred to as Amdahl's law. FIG. 1 is a graph illustrating the effects of Amdahl's law. As illustrated in FIG. 1, the speed up in latency of the execution of tasks from using additional parallel processors levels off even as additional parallel processors are added, due to the speedup being limited by the serial portion of the program.

(16) In some applications, a problem may consist of both a small time scale and a very large domain size. An example of this is direct numerical simulations of the Navier Stokes equation. Typically, these types of problems are never solved except on rare occasions on national super computers despite the unprecedented accuracy.

(17) In addition, in cases where the solution at a particular time interval needs to be recorded, then in most cases the computation will stop to download this timestep information, further adding to the time to solution for a given problem.

#### System Overview

(18) Embodiments are directed to a computer architecture specialized to solve differential equations that addresses the problems expounded above. FIG. 2 illustrates a high level diagram of a system for solving differential equations, in accordance with some embodiments. The system comprises an interface computer or circuit **202**, a dispatch computer or circuit **204**, and a plurality of solver units **206** (or solvers **206**). In some embodiments, the interface computer **202**, dispatch computer **204**, and solver units **206** are implemented on an application server **210**. While FIG. 2 illustrates a single application server **210**, it is understood that in other embodiments, the interface computer **202**, dispatch computer **204**, and solvers **206** may be implemented on multiple servers or devices, on a cloud server, etc.

(19) In some embodiments, the user accesses the application server **210** from a user device **212**, such as a PC, laptop, workstation, mobile device, etc. The user device **212** may access the application server **210** through a network **220** (e.g., the Internet). In other embodiments, the user device **212** may connect to the application server **210** via a direct line connection (e.g., a direct line connection to the interface computer **202**). In addition, although FIG. 1 only shows a single user device **212** connecting to the application server **210**, it is understood that in some embodiments, many user devices may concurrently connect to the application server **210** (e.g., via the network **220**).

(20) The user at the user device **212** may transmit to the application server **210** (e.g., through the network **220**) one or more problems involving differential equations to be solved. In some embodiments, the user device **212** transmits each problem in the form of the problem package, comprising at least a differential equation associated with the problem, and a domain. In some embodiments, the problem package further comprises a mesh (or particle domain) for the problem, one or more boundary conditions, initial conditions, flow conditions (such as density and viscosity), a solve type (e.g., 3D incompressible DNS Navier Stokes), and/or the like. The problem package may be sent to the interface computer **202** over the secured internet using a provided API of the interface computer **202**.

(21) The interface computer **202** receives the problem package, which is processed by the dispatch computer **204** and dispatched to the solvers **206**. The solvers **206** generate a solutions package comprising a solved domain that is transmitted back to the user device **212**. The solutions package may further comprise one or more averages, one or more solver metrics, one or more errors messages, etc.

(22) FIG. 3 illustrates a diagram of the interface computer **202**, dispatch computer **204**, and solver units **206**, in accordance with some embodiments. The interface computer **202** is networked to both

the user (e.g., the user device **212**) and the dispatch computer **204**. The interface computer **202** comprises a problem queue **302**, a solution queue **304**, and an error queue **306**. The interface computer **202** is configured to accept incoming problems to be solved by various interested parties (e.g., problem packages from one or more users at user devices **212**), and add the received problem packages **308** to the problem queue **302**. In some embodiments, the interface computer **202** may first check the received problem package for accuracy. For example, the interface computer **202** may, if the problem package specifies a time step size and is associated with certain types of differential equations, that the specified time step conforms with the Courant-Friedrichs-Lewy (CFL) convergence condition. In cases where the problem package specifies an unstructured mesh, the interface computer **202** may check if the specified mesh is well-formed. In some embodiments, the interface computer **202** may receive a problem package that comprises geometry information with initial and boundary conditions instead of a mesh, whereupon the interface computer may generate a mesh for the problem based upon the received geometry information and conditions.

(23) The interface computer **202** sends problem packages **308** to the dispatch computer **204** to be solved by one or more of the plurality of solvers **206**. In some embodiments, each problem package within the problem queue **302** may be assigned a priority level. The priority level for a problem package **302** may be based upon a provided indication within the problem package, the user from which the problem package was received, one or more parameters of the problem package (e.g., type of differential equation, size of domain, etc.), size of the problem package, an amount of time the problem package has been in the problem queue **302**, and/or any combination thereof.

(24) The interface computer **202** is further configured to receive solution information from the dispatch computer **204**. In some embodiments, the solution information is received in the form of one or more solution packages (e.g., as described above). In other embodiments, the interface computer **202** reformats the received solution information to form one or more solution packages. The interface computer **202** stores the one or more solution packages in the solution queue **304**, and transmits the solution packages from the solution queue **304** to their respective users (e.g., to the user devices **212** responsible for sending the problem package corresponding to the solution package).

(25) In some embodiments, the interface computer **202** receives error information from the dispatch computer **204**, corresponding to any errors encountered by the solvers **206** when solving the problem. In some embodiments, the dispatch computer **202** checks the fidelity of the results of the received solution information, and generates one or more errors if any issues are found (e.g., pressure, density, velocity, etc. parameters not being bounded). The determined errors may be stored in the error queue **306**, to be transmitted to corresponding users.

(26) The dispatch computer **204** is networked to the interface computer **202** and to one or multiple solver units **206**. The dispatch computer determines which solver **206** is the best to solve a given user problem at a given time. As illustrated in FIG. 3, the dispatch computer **204** may be in communication with a plurality of solvers **206** (e.g., solvers **206-1** through **206-n**). In some embodiments, the dispatch computer **204** monitors an availability of the solvers **206** (e.g., a capacity of each solver to process additional problems) and the problem queue **302** of the interface computer **202**, in order to determine which problem packages **308** should be processed by which of the solvers **206**.

(27) The solvers **206** are the workhorses of the system, and are configured to generate solutions to the various problems that come to the system. The solvers may be of different types. For example, each of the solvers **206** may be optimized for one or more specific applications, such as fluid dynamics, molecular dynamics, electronic structure, etc. In some embodiments, each solver **206** may also be optimized to solve domains of different sizes. The various sizes may help optimize the use of the hardware by allocating larger problems to the larger solvers and smaller problems to the smaller solvers.

(28) FIG. 4 is a diagram illustrating components of a solver, in accordance with some

embodiments. The solver **400** illustrated in FIG. 4 may correspond to one of the solvers **206** illustrated in FIGS. 2 and 3. The solver **400** comprises a coordinator computer or circuit **402**, multiple compute units (referred to as Differential Equation Accelerator (DEA) circuits, or DEAs) **404**, one or more DEA—Coordinator interconnects **406**, and one or more DEA—DEA interconnects **408**.

(29) The coordinator computer **402** (or coordinator **402**) is connected to the dispatch computer (e.g., dispatch computer **204**) on one side and to multiple DEAs **404** on the other. The coordinator **402** is responsible for coordinating the various aspects of the DEA when solving a user problem. For example, the coordinator **402** may, in response to receiving a problem package, divide the domain of the problem into a plurality of subdomains, and assigns each subdomain to a respective DEA **404**. The coordinator **402** may synchronize the DEAs **404** and initiates solving operations by the DEAs **404**. The coordinator **402** further downloads results from each of the DEAs **404**.

(30) The solver **400** comprises a plurality of DEAs **404**. Each DEA **404** is a circuit configured to receive a subdomain of a problem, and generate solution data for the received subdomain. The solution data may correspond to the subdomain processed over one or more time steps, based upon the differential equation associated with the problem. For the purposes of discussion, herein, the DEAs will be referred to as processing the subdomain data over one or more time steps, although it is understood that, depending on the type of problem to be solved, iterations other than time steps may be used. In some embodiments, the DEAs within each solver may be of the same type (e.g., having the same memory capacity, contain the same types of systolic arrays, etc.). Solvers of different sizes or configured to process different types of differential equation problems will have different DEAs.

(31) The coordinator computer **402** and the DEAs **404** are connected via DEA—Coordinator interconnects **406** and DEA—DEA interconnects **408**, allowing for the coordinator **402** to manage operations of the DEAs **404**, and for the DEAs **404** to share stored domain information with each other (discussed in greater detail below).

(32) The DEA—Coordinator interconnects **406** and DEA—DEA interconnects **408** may be implemented as cabling connecting the coordinator **402** to the DEAs **404**, and the DEAs **404** to each other, respectively. In some embodiments the interconnects **406** and **408** may be implemented using PCI Express cables (e.g., PCIe v4.0). The number of interconnects between the DEAs **404** may be contingent on how the domain is sliced up across the DEAs in that solver e.g., based on a partitioning scheme of the solver for partitioning received domains. For example, if the solver is configured to slice the domain into pyramids, then the number of interconnects may be smaller compared to if the solver were configured to slice up received domains into higher order polygons. In some embodiments, if the number of DEAs **404** is large, then it may be hard to physically connect all the DEAs **404** onto one coordinator **402**. In such cases, relays (not shown) can be used to bunch up some of the cabling.

(33) FIG. 5 is a flowchart of a process for solving a differential equation, in accordance with some embodiments. The process **500** of FIG. 5 may be performed by the interface computer **202**, dispatch computer **204**, and plurality of solvers **206** illustrated in FIGS. 2 and 3.

(34) The interface computer receives **505** a problem package from a user. In some embodiments, the problem package is formulated by the user at a user device (e.g., the user device **212**), and may include a mesh or particle domain associated with the problem, one or more boundary conditions, one or more initial conditions, one or more flow conditions (e.g., density and/or viscosity), a solve type (e.g., indicating a type of differential equation to be solved, such as a 3D incompressible DNS Navier-Stokes problem), a number of time steps and/or convergence criteria, a numerical method, and/or the like. For example, a particular problem package may specify that the problem relates to 3D flow around an airfoil using DNS incompressible Navier-Stokes equations, wherein the fluid is air at sea level, and the domain is a structured mesh. In some embodiments, the user at the user device sends the problem package to the interface computer using a provided API.

(35) In some embodiments, the interface computer checks **510** the received problem package for errors, and, if there are no errors, enqueues the problem package in a problems queue (e.g., problem queue **302**). In some embodiments, if errors are detected, the interface computer notifies the user and ejects the problem package. In some embodiments, each problem package enqueued in the problem queue may be associated with a priority level. The priority level may be based upon a default value, a priority level associated with the user submitting the problem package, an amount of time the problem package has been in the problem queue without being dispatched for solving, or any combination thereof.

(36) The dispatch computer dispatches **515** problem packages from the problem queue to one of the several solvers attached to it. In some embodiments, the dispatch computer is used to optimize the use of the solvers by choosing between the various types and sizes of the solvers attached to it. The size of a solver may be based upon a number of DEAs of the solver, an amount of memory space within each DEA, a size of the systolic arrays within each DEA of the solver, and/or a type of differential equation the systolic arrays of the DEAs are configured to solve. Additional details regarding how the dispatch computers determines how to dispatch problems to solvers is described in relation to FIG. **6** below.

(37) The solver that receives the problem package is used to compute a solution for the problem package. To compute a solution to the problem, the coordinator of the solver analyzes the problem package to determine a domain (e.g., based on the spatial domain, boundary conditions, and/or initial conditions indicated by the problem package). In some embodiments, the coordinator may, if a mesh was not included in the problem package, generate a mesh based on the determined domain. The coordinator generates **520** subdomains from the domain, and assigns the subdomains to the DEAs of the solver. In some embodiments, the subdomains are determined by slicing the domain based upon some virtual geometry to fit the memory spaces of the DEAs of the solver. Each of the DEAs assigned a subdomain receives its respective subdomain data, as well as additional data such as fluid viscosity and density needed to solve the problem.

(38) In some embodiments, the DEAs of the solver are configured to solve their respective subdomains concurrently. This allows for the DEAs to share subdomain data with each other during the processing for a given time step as needed, since all DEAs will be processing the same time step over a given time period. The coordinator initiates **525** solving when all DEAs of the solver having an assigned subdomain are ready. In some embodiments, each DEA signals to the coordinator when it has received its assigned subdomain data and is ready to begin solving, whereupon the coordinator issues a signal to begin solving when all DEAs having assigned subdomains have indicated that they are ready.

(39) The DEAs each solve the problem for their respective subdomain by processing the subdomain over a plurality of time steps based upon the differential equation of the problem. The coordinator receives **530** the solved subdomains from the DEAs, and assembles the solved subdomains from the various DEAs to generate a solved domain corresponding to the original domain of the problem.

(40) The solver sends **535** the solution (e.g., the solved domain) to the dispatch computer, which performs one or more quality control functions on the received solved domain. In some embodiments, the dispatching computer checks the fidelity of the results. For example, Table 1 below shows some of the parameters used in Navier Stokes equations that may be checked by the dispatch computer. In some embodiments, the parameters checked may be based upon a type of equation associated with the problem package.

(41) TABLE-US-00001 TABLE 1 Parameter (all values in domain) Check Pressure, density, velocity, Bounded temperature, energy Temperature Not less than 0 K Density, pressure Not less than 0

(42) If no errors are found, the dispatch computer may enqueue the solved domain on the solution queue of the interface computer, from where the solved domain can be transmitted **540** back to the



user. The solved domain may be transmitted as part of a solution package, which may include the solved domain as well as averages, solver metrics, and/or other information.

(43) If errors are found by the dispatch computer in the solved domain, the dispatch computer may send the problem to be solved by a different solver. Otherwise (or after at least a threshold number of different solvers have attempted to solve the problem), the error information may be enqueued in the error queue of the interface computer to be sent back to the user. In some embodiments, the user may receive, from the interface computer, the solved domain data along with any corresponding error information.

(44) FIG. 6 is a flowchart illustrating a process for dispatching problems to solvers, in accordance with some embodiments. As described above, the dispatch computer is used to optimize the use of the solvers by choosing between the various types and sizes of the solvers attached to it. The dispatch computer may dispatch problems to the solvers, based upon the priority of the enqueued problems, and the availability of solvers.

(45) The dispatch computer monitors **605** each of a plurality of solvers. For example, the dispatch computer may maintain information indicating which solvers are currently processing which problems, and which solvers are available to take on new problems.

(46) The dispatch computer determines **610** if a solver of the plurality of solvers is available to solve a new problem. In some embodiments, when a solver finishes solving a problem, the solver notifies the dispatch computer that it is done and ready to receive a new problem. If no solvers are currently available, the dispatch computer continues to monitor the solvers.

(47) If a solver is available, the dispatch computer makes a determination on how to utilize the solvers to solve the problems currently enqueued in the problem queue. As discussed above, each enqueued problem package may be associated with a corresponding priority value. In some embodiments, the dispatch computer determines **615** if there is currently a problem package in the problem queue having a highest level of priority (or a priority level exceeding a threshold value).

(48) If a problem package having the highest priority level exists, the dispatch computer determines **620** if the available solver is able to fit the problem. In some embodiments, each solver of the plurality of solvers may be optimized to solve domains of different sizes. In addition, different solvers may be configured to solve different types of differential equations. For example, if the available solver is configured to solve problems associated with the differential equation of the problem with domain sizes smaller than that of the identified problem package, the solver may not be able to solve the problem package. If the available solver is able to fit the problem (e.g., is configured to solve differential equations of the same type as those associated with the problem package and to solve domains at least the size of that of the problem package), the dispatch computer dispatches **625** the identified problem package having high priority to the solver.

(49) On the other hand, if the available solver does not fit the problem, or if there is no queued problem having a high enough priority level, the dispatch computer may attempt to determine a scheme for efficiently distributing the queued problems to available solvers. In some embodiments, the dispatch computer determines **630** an estimated finish time for each monitored solver expected to finish solving its current problem within a threshold period of time. In some embodiments, the dispatch computer determines the estimated finish time for a solver by determining an estimated solve time for the solver (based on the known domain size of the problem package currently being solved by the solver, the type of differential equation, the size of the solver, etc.) and compares the estimated solve time to how much time the solver has already spent processing the problem. In some embodiments, the dispatch computer may query the solver to determine a current status of the solver (e.g., based upon a current time-step being processed by the solver) from which an estimated finish time can be determined. In some embodiments, the dispatch computer determines an initial estimated solve time for the solver when dispatching a problem package to the solver, and later may check or adjust the estimate based upon one or more queries to the solver.

(50) The dispatch computer determines **635** one or more permutations to generate one or more

solver scenarios. Each solver scenario corresponds to a different arrangement of how problems in the problem queue may be able to fit in the available solvers and solvers expected to become available with a predetermined time period (e.g., next n minutes). The dispatch computer analyzes the possible scenarios, each corresponding to a different permutation of problems assigned to solvers, and determines how well each scenario utilizes the available solvers/DEAs. In some embodiments, the dispatch computer determines a utilization score for each scenario, based upon a number or percentage of DEAs within each solver that will be utilized for solving, a number of problems able to be dispatched, an estimated finish time of the utilized solvers, or some combination thereof. For example, if a particular solver is able solve problems having a large domain, assigning a problem having a small domain to that solver may involve only a portion of the DEAs of the solver being assigned subdomains associated with the problem, potentially resulting in a lower DEA utilization of the solver. In some embodiments, the utilization score may further be based upon a priority level of problems dispatched to the solvers (e.g., a first scenario where more higher priority problems are dispatched to solvers may have a higher utilization score than a second scenario, even if the second scenario may result in the dispatch of more problems). (51) The dispatch computer selects **640** a scenario having a best solver/DEA utilization (e.g., highest utilization score), and sends queued problem packages to solvers in accordance with the selected scenario to be solved.

(52) In some embodiments, a selected scenario may not be able to assign each queued problem to a solver, such as if the number of queued problems exceeds a number of available solvers. In some embodiments, the dispatch computer determines **650** if there are any problems overlooked by the selected scenario, and as such remain in the queue. If there are overlooked problems in the queue, the dispatch computer may cause **655** a priority level of the overlooked problems to be increased. For example, the dispatch computer may increment the priority level of each overlooked problem by a set amount, or in accordance with a predetermined function. The dispatch computer may then continue to monitor the solvers. Using this scheme, the dispatch computer is able to dispatch problems to solvers in such a way that will maximum solver/DEA utilization, while also ensuring that problems do not remained enqueued for too long (e.g., by only determining dispatch permutation scenarios if no currently queued problem has reached a threshold priority level).

#### Solver DEA Structure

(53) FIG. 7 illustrates a layout of a DEA **404**, in accordance with some embodiments. The DEA comprises a coordinator-DEA interconnect and controller **702**, which is a special circuit and interconnect that manages the data and control signals going back and forth between the DEA unit and the coordinator **402**. For example, the DEA may receive problem and subdomain data from the coordinator **402** via the coordinator-DEA interconnect and controller **702**. The DEA may also receive instructions from the coordinator **402** (e.g., synchronization instructions to synchronize with other DEAs of the solver, instructions to begin solving, etc.). In addition, the DEA may transmit generated solution information corresponding to the received problem and subdomain back to the coordinator through the coordinator-DEA interconnect and controller **702**.

(54) The external memory interconnect & controller **704** is a special circuit and interconnect that manages data and control signals between the various DEA units. For example, as will be discussed in greater detail below, in some embodiments, the DEA may require information relating to other subdomains being processed by other DEAs of the solver. As such, the DEA may receive additional subdomain data from other DEAs via the external memory interconnect & controller **704**.

(55) The control module **706** manages the overall functioning of the DEA unit. In some embodiments, the control module **706** is a processor that processes received subdomain data, determines and stores parameters associated with the problem subdomain (e.g., in the parameters storage **708**), and manages solving of the problem subdomain (e.g., using the processing element **714**) over a plurality of time steps.

(56) The parameters storage **708** is configured to store local variables used during the solving of

differential equations. In some embodiments, the parameters storage **708** is implemented as an SRAM. The stored local variables may include any type of variable expected to be highly used during solving of the problem assigned to the DEA that are not expected to change during the solving, such as subdomain data, solve type, and one or more constants to be used during the solving of the subdomain (e.g., fluid density, viscosity, etc.).

(57) The memory **710** is used to store the problem to be solved. In some embodiments, the memory **710** each DEA is divided into three subunits (e.g., first memory unit **710-1**, second memory unit **710-2**, and third memory unit **710-3**). In some embodiments, the memory units **710-1** through **710-3** are implemented as part of the same memory. In other embodiments, the memory units **710-1** through **710-3** are implemented as two or more separate memory chips.

(58) In some embodiments, first and second memory units **710-1** and **710-2** are used in general solving of the differential equation, while the third memory unit **710-3** may be used when the DEA needs to send data back to the coordinator (e.g., via the coordinator-DEA interconnect and controller **702**). In some embodiments, access to the memory units **710-1** to **710-3** is managed by the internal memory controller **712**. For example, the internal memory controller **712** may receive instructions from the control module **706** to retrieve data between the first and second memory units **710-1** and **710-2** and the processing element **714**, move processed data to the third memory unit **710-3** in preparation for transmission to the coordinator of the solver, and/or the like.

(59) The processing element **714** is configured to receive problem data (e.g., from the first or second memory units **710-1** and **710-2**) and to solve the received problem data using one or more systolic arrays. In some embodiments, the processing element **714** comprises one or more gatekeeper circuits **716** (also “gatekeepers **716**”) and a plurality of systolic array circuits **718** (“systolic arrays **718**”). The gatekeepers **716** are circuits that divert data from memory (e.g., from the first or second memory units **710-1** and **710-2** via the internal memory controller **712**) to the systolic arrays **718** and vice versa, depending on which equation is solved. For example, the gatekeeper **716** may receive information indicating a type of differential equation to be solved from the parameters storage **708** where solver parameters are kept, and select which systolic array **718** to use to process problem data received from the first memory unit **710-1** or the second memory unit **710-2**.

(60) The systolic arrays **718** each comprise hardware configured to solve a particular type of partial differential equation (PDE). Each systolic array **718** comprises a network of elements (e.g., processing elements and storage elements) configured to process a plurality of nodes of a received (sub-)domain in parallel. Each processing element of the systolic array is configured to compute a partial result as a function of the data received from its upstream elements, and to pass results to downstream elements. In some embodiments, the systolic arrays **718** comprise at least one systolic array for each type of PDE that the DEA is designed to solve (e.g., PDE1 through PDEn). For example, a systolic array may be configured to solve 1-D differential equations such as linear convection, non-linear convection, diffusion, Burger's equation, Laplace equation, Poisson equation, Euler's equation, Navier stokes simulations, etc. In some embodiments, a systolic array may be configured to solve a multi-dimensional differential equation. In some embodiments, depending on the similarity of the equation, different PDEs may be solved on the same systolic array with minor changes to the calculation made by gates of the systolic array based upon parameters provided by the parameters storage **508**.

(61) The DEA is configured to solve a subdomain of the problem sent to it by the coordinator. FIG. **8** is a flowchart of a process for using a DEA to solve a subdomain, in accordance with some embodiments. The DEA receives **805** subdomain data from the coordinator of the solver. The received subdomain data may be copied into the third memory unit **710-3**. In addition, the DEA may receive other data of the problem package, such as flow conditions, solve type, and constants such as fluid density and viscosity, which is copied into the parameters module **708**. In some embodiments, each DEA may also receive external domain data part of the initial mesh,

corresponding to portions of subdomains adjacent to the DEAs assigned subdomain, for use in determining the first time step for the subdomain.

(62) Once all data has been disseminated by the coordinator to the DEAs of the solver, the DEA is synchronized **810** with the other DEAs of the solver by the coordinator. When all DEAs of the solver having an assigned subdomain are synced, the coordinator initiates solving, during which each DEA processes its assigned subdomain over a plurality of time steps. In some embodiments, the DEA begins solving **815** by processing the first time step of the whole subdomain stored in the third memory unit **710-3** using the processing element **714**, storing the results of the processing into the first memory unit **710-1**. The process may take 1 or more clock cycles. In some embodiments, memory sharing between DEAs may not need to be performed at this point, as external domain data corresponding to data from other subdomains needed to process the first time step for the subdomain may have been received from the initial mesh.

(63) The DEAs of the solver are synchronized to concurrently perform each time step. In some embodiments, a DEA of the solver may communicate with one or more adjacent DEAs regarding their current status, allowing for the DEAs of the solver having assigned subdomains to begin processing of each time step synchronously. During each time step, the DEAs share **820** parts of subdomain stored in the first memory unit **710-1** as needed (e.g., subdomain data near a boundary of the another subdomain, which is needed by another DEA assigned the another subdomain for processing a next time step). In addition, each DEA processes **825** its respective subdomain through the processing element **714** to determine the next time step for the subdomain. In some embodiments, the processing element **714** receives the subdomain data from the first memory unit **710-1**, and selects an appropriate systolic array **718** to be used for processing, using stored parameters from the parameters storage **708**. The results of the processing are stored in the second memory unit **710-2**. Although FIG. 8 illustrates **820** and **825** as separate steps, it is understood that these two steps may be performed concurrently.

(64) The DEAs may repeat time stepping over a plurality of cycles (steps **820** and **825**). Over each time step, the processing element receives the subdomain data from the first or second memory unit **710-1** or **710b**, selects a systolic array for processing the subdomain data, and stores the processed subdomain data into the opposite memory unit (e.g., from the first memory unit **710-1** to the second memory unit **710-2**, or vice versa). In addition, memory sharing with other DEAs may be performed concurrently. In some embodiments, the same systolic array may be used for each time step. In other embodiments, different systolic arrays may be selected, based upon the problem being solved. For example, when solving a combustion problem, a first pass may comprise one or more time steps in which a systolic array for solving fluid dynamics is selected, and a second pass may comprise one or more time steps using a systolic array for solving for the chemistry.

(65) In some embodiments, if a data extraction for a particular time step is needed, then the processing element **714** may also output **830** the processed data to the third memory unit **710-3** along with to first or second memory unit **710-1** or **710b**. The DEA may then inform the coordinator of the solver to download the time step data from the third memory unit **710-3**. In addition, the DEA may concurrently continue to solve between the first and second memory units **710-1** and **710b**, since the data download could last more than one clock cycle. In some embodiments, the DEA is configured to output its data to the third memory unit **710-3** for download by coordinator as “snapshots” at predetermined intervals (e.g., every predetermined number of time-steps) or in accordance with a predetermined function. In other embodiments, these snapshots may be taken dynamically. The snapshot data may be used to analyze how the solution of the problem package evolves over time, and/or to perform accuracy checks (e.g., verify that momentum or energy are conserved).

(66) In some embodiments, a number of time steps to be solved may be explicitly indicated as part of the problem package. In other embodiments, the problem may be implicit, in which the solver solves until a specified parameter reaches a predetermined value. For example, the solver, at each

time step, may check the root mean square of the velocities and stop solving once it has reached a certain critical value. In some cases, a maximum number of steps may be specified, in case the aggregate critical value is not reached.

(67) At the completion of solving, the final processed data may be output **835** by the processing element **714** to the third memory unit **710-3**. The DEA then informs the coordinator to download the solved data for the subdomain from the third memory unit **710-3**. The coordinator, upon downloading the solved subdomain data from each of the plurality of DEAs, may assemble the solved subdomain data into a solutions package for the domain.

(68) By dividing the memory of the DEA into first and second memory units **710-1** and **710-2**, the DEA ensures that processed data for each time step can be written to memory without disturbing the original pre-time step data until processing of the entire subdomain is completed. In addition, the third memory unit **710-3** allows for data to be extracted and sent to the coordinator without interrupting time step processing of the subdomain data.

(69) The language used in the specification has been principally selected for readability and instructional purposes, and it may not have been selected to delineate or circumscribe the inventive subject matter. It is therefore intended that the scope of the patent rights be limited not by this detailed description, but rather by any claims that issue on an application based hereon.

Accordingly, the disclosure of the embodiments is intended to be illustrative, but not limiting, of the scope of the patent rights, which is set forth in the following claims.

## Claims

1. A system for solving differential equations, comprising: a plurality of solver units, each comprising a plurality of differential equation accelerators (DEAs); a dispatch circuit configured to identify one or more problem packages, each problem package indicating at least a differential equation to be solved and a domain, and dispatch a problem package of the identified problem packages to a selected solver unit of the plurality of solver units, wherein the selected solver unit is configured to, responsive to receiving the problem package: divide the domain of the problem package into a plurality of subdomains; assign each of the plurality of subdomains to a respective DEA of the solver unit, each DEA comprising hardware configured for incrementally solving its assigned subdomain based upon the differential equation over a plurality of steps; in response to determining that all DEAs of the solver unit having an assigned subdomain have received data corresponding to their respective assigned subdomain, initiate solving by the DEAs, such that each of the DEAs having an assigned subdomain solves for a particular step concurrently; receive solved subdomains from the respective DEAs having assigned subdomains; and assemble the received solved subdomains into a solved domain.

2. The system of claim 1, where the dispatch circuit is configured to: monitor the plurality of solver units; responsive to a determination that a solver unit of the plurality of solver units is available and that the problem package one of the one or more problem packages has a priority level of at least a threshold level, dispatch the problem package to the solver unit if the solver unit fits the problem package.

3. The system of claim 1, wherein the dispatch circuit is further configured to: responsive to determining that the one or more identified problem packages does not include a problem package having a priority level of at least a threshold level, determine estimated finish times for solver units of the plurality of solver units that are within a predetermined time period; based upon the estimated finish times, determine one or more scenarios, each corresponding to a permutation of problem packages of the one or more problem packages assigned to solver units of the plurality of solver units; determine, for each scenario, a utilization metric based at least in part on a level of utilization of DEAs of the plurality of solver units under the scenario; select a scenario of the one or more scenarios, based upon the determined utilization metric for each scenario; and dispatch

problem packages of the one or more problem packages to solver units of the plurality of solver units in accordance with the selected scenario.

4. The system of claim 3, wherein the dispatch circuit is further configured to: determine that the one or more problem packages includes one or more additional problem packages not assigned to a solver unit of the plurality of solver units in accordance with the selected scenario; increase a priority level of each of the one or more additional problem packages.

5. The system of claim 3, wherein the utilization metric for a scenario of the one or more scenarios is further based upon a number of problem packages assigned to solver units under the scenario and an estimated finish time associated with the solver units assigned problem packages under the scenario.

6. The system of claim 3, wherein the level of utilization of DEAs of the plurality of solver units for a scenario of the one or more scenarios is based upon a percentage of DEAs of a solver unit of the plurality of solver units that would be utilized to solve a problem package assigned to the solver unit under the scenario.

7. The system of claim 3, wherein the dispatch circuit is configured to determine estimated finish times for solver units of the plurality of solver units that are within a predetermined time period comprises by determining, for a solver unit of the plurality of solver units, an estimated solve time for the solver unit based a problem package currently being solved by the solver unit, and an amount of time already spent by the solver unit in processing the problem package.

8. The system of claim 1, wherein DEAs of the selected solver unit having respective assigned subdomains are further configured to communicate with each other to synchronize solving for each step of the plurality of steps.

9. A system for solving differential equations, comprising: a plurality of solver units; a dispatch circuit configured to dispatch problem packages of a plurality of problem packages to solver units of the plurality of solver units, wherein each problem package indicates at least a respective differential equation to be solved and a domain, wherein the dispatch circuit is configured to: determine estimated finish times for solver units of the plurality of solver units; based upon the estimated finish times, determine one or more scenarios, each corresponding to a permutation of problem packages of the plurality of problem packages assigned to solver units of the plurality of solver units; select a scenario of the one or more scenarios, based upon a determined utilization metric of the plurality of solver units for each scenario; and dispatch problem packages of the plurality of problem packages to solver units of the plurality of solver units in accordance with the selected scenario.

10. The system of claim 9, wherein the dispatch circuit is configured to determine the one or more scenarios based upon solver units of the plurality of solver units having estimated finish times within a predetermined time period.

11. The system of claim 9, wherein the dispatch circuit is further configured to, responsive to determining that the plurality of problem packages includes one or more additional problem packages not assigned to a solver unit of the plurality of solver units in accordance with the selected scenario, increase a priority level of each of the one or more additional problem packages.

12. The system of claim 9, wherein the dispatch circuit is configured to determine the utilization metric for the scenario based upon a number of problem packages assigned to solver units under the scenario and an estimated finish time associated with the solver units assigned problem packages under the scenario.

13. The system of claim 9, wherein each of the plurality of solver units comprises a respective plurality of differential equation accelerators (DEAs), each DEA comprising hardware configured for incrementally solving an assigned subdomain based upon a differential equation over a plurality of steps.

14. The system of claim 13, wherein the dispatch circuit is configured to determine the utilization metric for the scenario based upon a level of utilization of DEAs of the plurality of solvers under

the scenario.

15. A non-transitory computer readable medium comprising stored instructions that, when executed by at least one processor, configure the at least one processor to dispatch problem packages of a plurality of problem packages to solver units of a plurality of solver units, based upon an availability of the plurality of solver units, wherein each problem package indicates at least a respective differential equation to be solved and a domain, by: determining estimated finish times for solver units of the plurality of solver units; based upon the estimated finish times, determining one or more scenarios, each corresponding to a permutation of problem packages of the plurality of problem packages assigned to solver units of the plurality of solver units; selecting a scenario of the one or more scenarios, based upon a determined utilization metric of the plurality of solver units for each scenario; and dispatching problem packages of the plurality of problem packages to solver units of the plurality of solver units in accordance with the selected scenario.

16. The non-transitory computer readable medium of claim 15, wherein the stored instructions, when executed by the at least one processor, further configure the at least one processor to determine the one or more scenarios based upon solver units of the plurality of solver units having estimated finish times within a predetermined time period.

17. The non-transitory computer readable medium of claim 15, wherein the stored instructions, when executed by the at least one processor, further configure the at least one processor to, responsive to determining that the plurality of problem packages includes one or more additional problem packages not assigned to a solver unit of the plurality of solver units in accordance with the selected scenario, increase a priority level of each of the one or more additional problem packages.

18. The non-transitory computer readable medium of claim 15, wherein the stored instructions, when executed by the at least one processor, further configure the at least one processor to determine the utilization metric for the scenario based upon a number of problem packages assigned to solver units under the scenario and an estimated finish time associated with the solver units assigned problem packages under the scenario.

19. The non-transitory computer readable medium of claim 15, wherein each of the plurality of solver units comprises a respective plurality of differential equation accelerators (DEAs), each DEA comprising hardware configured for incrementally solving an assigned subdomain based upon a differential equation over a plurality of steps.

20. The non-transitory computer readable medium of claim 19, wherein the stored instructions, when executed by the at least one processor, further configure the at least one processor to determine the utilization metric for the scenario based upon a level of utilization of DEAs of the plurality of solvers under the scenario.

---