

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12393643
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Soryal; Joseph et al.

---

### Audio attack mitigation system with alarm indication

---

#### Abstract

Concepts and technologies disclosed herein are directed to image classification attack mitigation. According to one aspect of the concepts and technologies disclosed herein, a system can obtain an original image and reduce a resolution of the original image to create a reduced resolution image. The system can classify the reduced resolution image and output a first classification. The system also can classify the original image via deep learning image classification and output a second classification. The system can compare the first classification and the second classification. In response to determining that the first classification and the second classification match, the system can output the second classification of the original image. In response to determining that the first classification and the second classification do not match, the system can output the first classification of the original image.

---

<b>Inventors:</b>	Soryal; Joseph (Glendale, NY), Reid; Dylan C. (Atlanta, GA)
<b>Applicant:</b>	SimpliSafe, Inc. (Boston, MA)
<b>Family ID:</b>	1000008767465
<b>Assignee:</b>	SimpliSafe, Inc. (Boston, MA)
<b>Appl. No.:</b>	18/623085
<b>Filed:</b>	April 01, 2024

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20240248958 A1	Jul. 25, 2024

#### Related U.S. Application Data

continuation parent-doc US 18117622 20230306 US 11947630 child-doc US 18623085  
continuation parent-doc US 17218635 20210331 US 11599754 20230307 child-doc US 18117622

---

## Publication Classification

**Int. Cl.:** G06F18/22 (20230101); G06F18/20 (20230101); G06F18/21 (20230101); G06V30/24 (20220101)

**U.S. Cl.:**

**CPC** G06F18/217 (20230101); G06F18/22 (20230101); G06F18/285 (20230101); G06V30/2504 (20220101);

## Field of Classification Search

**CPC:** G06F (18/217); G06F (18/22); G06F (18/285); G06V (30/2504)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2006/0114058	12/2005	Chen	330/10	H03F 3/181
2007/0046503	12/2006	van Baarsen	341/50	H03M 5/22
2012/0109646	12/2011	Bak	704/229	G10L 15/065
2013/0268956	12/2012	Recco	N/A	N/A
2019/0087677	12/2018	Wolf et al.	N/A	N/A
2019/0205649	12/2018	Ananthanarayanan et al.	N/A	N/A
2020/0019128	12/2019	Brooks et al.	N/A	N/A
2021/0092280	12/2020	Nishimura et al.	N/A	N/A
2021/0343291	12/2020	Agarwal	N/A	G10L 15/22
2022/0013118	12/2021	Hwang	N/A	G10L 19/018
2022/0122614	12/2021	Pelecanos et al.	N/A	N/A
2022/0138947	12/2021	Kuo	N/A	N/A
2022/0148598	12/2021	Krishnamoorthy	N/A	G10L 15/20

### OTHER PUBLICATIONS

U.S. Notice of Allowance dated Oct. 5, 2022 in U.S. Appl. No. 17/218,635. cited by applicant  
U.S. Office Action dated Jun. 24, 2022 in U.S. Appl. No. 17/218,635. cited by applicant

---

*Primary Examiner:* Bhattacharya; Sam

*Attorney, Agent or Firm:* Pierce Atwood LLP

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 18/117,622, filed Mar. 6, 2023, which is a continuation of U.S. patent application Ser. No. 17/218,635, filed Mar. 31, 2021 (now U.S. Pat. No. 11,599,754). All sections of the aforementioned application(s) and patent(s) are incorporated herein by reference in their entirety.

## BACKGROUND

(1) The field of computer vision utilizes artificial neural networks inspired by the organization of neurons in the visual cortex of the human brain. Convolutional neural networks (“CNNs”) are the most widely used artificial neural networks for analyzing and classifying images. CNNs use deep learning algorithms to assign weights to various aspects or objects depicted in an image to differentiate the image from other images and to assign a classification to the image. Image classification has become the most prevalent use case for artificial intelligence. As with any prevalent technology, attackers will find ways to exploit the technology for malicious purposes.

(2) Attackers can breach image classification systems and insert malicious pixels into images in an image feed to trick the artificial intelligence to misinterpret an image and provide an incorrect classification. For example, attackers may want to cause an image classification system to interpret an image of an animal as a gun. In particular, attackers can exploit the process of elimination that the image classification system uses when estimating which label to apply to an image.

Characteristics can be extracted from the image that is most likely to be classified as a first thing, and then applied imperceptibly to images of a second thing so that images of the first thing become classified as the second thing. The mathematics that power the elimination process allow an attacker to systematically push a poisoned image towards a target classification.

## SUMMARY

(3) Concepts and technologies disclosed herein are directed to image classification attack mitigation. According to one aspect of the concepts and technologies disclosed herein, a system can obtain an original image and reduce a resolution of the original image to create a reduced resolution image. The system can classify the reduced resolution image and output a first classification. The system also can classify the original image via deep learning image classification and output a second classification. The system can compare the first classification and the second classification. In response to determining that the first classification and the second classification match, the system can output the second classification of the original image. In response to determining that the first classification and the second classification do not match, the system can output the first classification of the original image.

(4) The system can attempt to reconstruct the original image from the first classification. The system can compare a reconstructed image to the original image. In response to determining that the reconstructed image matches the original image, the system can determine that the original image was accurately processed. In response to determining that the reconstructed image does not match the original image, the system can adjust the resolution of the original image and repeat classification.

(5) In some embodiments, the system can classify the reduced resolution image, at least in part, by performing an elimination operation using color as a primary classifier and shape as a secondary classifier. In some embodiments, the system can slice the reduced resolution image into individual items and search for common coexisting items.

(6) In some embodiments, the system can classify the reduced resolution image based upon other factors. For example, the system can perform an environment context awareness check on the reduced resolution image, a situational context awareness check on the reduced resolution image, a textual relationship check on the reduced resolution image, an audible relationship check on the reduced resolution image, a user profile biasing on the reduced resolution image, and/or a relative dimension and mathematical ratio analysis on the reduced resolution image.

(7) It should be appreciated that the above-described subject matter may be implemented as a computer-controlled apparatus, a computer process, a computing system, or as an article of manufacture such as a computer-readable storage medium. These and various other features will be apparent from a reading of the following Detailed Description and a review of the associated drawings.

(8) This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended that this Summary be used to limit the scope of the claimed subject matter. Furthermore, the claimed subject matter is not limited to implementations that solve any or all disadvantages noted in any part of this disclosure.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram illustrating an image classification attack mitigation system in which aspects of the concepts and technologies disclosed herein can be implemented.
- (2) FIG. 2 is a flow diagram illustrating aspects of a method for mitigating image classification attacks, according to an illustrative embodiment.
- (3) FIG. 3 is a flow diagram illustrating aspects of a method for classifying an image to determine an attack mitigation image classification, according to an illustrative embodiment.
- (4) FIG. 4 is a flow diagram illustrating aspects of a method for reconstructing the image from the attack mitigation image classification, according to an illustrative embodiment.
- (5) FIG. 5 is a flow diagram illustrating aspects of a method for performing image processing to classify the image and compare the classification to the attack mitigation image classification, according to an illustrative embodiment.
- (6) FIG. 6 is a block diagram illustrating an audio attack mitigation system in which aspects of the concepts and technologies disclosed herein can be implemented.
- (7) FIG. 7 is a flow diagram illustrating aspects of a method for comparing text generated from multiple audio samples, according to an illustrative embodiment.
- (8) FIG. 8 is a flow diagram illustrating aspects of a method for comparing text generated from enhanced audio and compressed audio, according to an illustrative embodiment.
- (9) FIG. 9 is a flow diagram illustrating aspects of a method for comparing text outputs, according to an illustrative embodiment.
- (10) FIG. 10 is a block diagram illustrating an example computer system capable of implementing aspects of the embodiments presented herein.
- (11) FIG. 11 is a block diagram illustrating an example containerized cloud architecture and components thereof capable of implementing aspects of the embodiments presented herein.
- (12) FIG. 12 is a block diagram illustrating an example virtualized cloud architecture and components thereof capable of implementing aspects of the embodiments presented herein.
- (13) FIG. 13 is a diagram illustrating a machine learning system, according to an illustrative embodiment.
- (14) FIG. 14 is a diagram illustrating a network, according to an illustrative embodiment.

### DETAILED DESCRIPTION

(15) While the subject matter described herein may be presented, at times, in the general context of program modules that execute in conjunction with the execution of an operating system and application programs on a computer system, those skilled in the art will recognize that other implementations may be performed in combination with other types of program modules.

Generally, program modules include routines, programs, components, data structures, computer-executable instructions, and/or other types of structures that perform particular tasks or implement particular abstract data types. Moreover, those skilled in the art will appreciate that the subject matter described herein may be practiced with other computer systems, including hand-held devices, mobile devices, wireless devices, multiprocessor systems, distributed computing systems, microprocessor-based or programmable consumer electronics, minicomputers, mainframe computers, routers, switches, other computing devices described herein, and the like.

(16) Referring now FIG. 1, a block diagram illustrating an image classification attack mitigation (“ICAM”) system **100** in which aspects of the concepts and technologies disclosed herein can be implemented will be described. The ICAM system **100** can be implemented, at least in part, in a computer system, such as an example computer system **1000** that is illustrated and described with reference to FIG. 10. The ICAM system **100** alternatively can be implemented, at least in part, in a containerized architecture, such as an example containerized cloud architecture **1100** that is illustrated and described herein with reference to FIG. 11. The ICAM system **100** can be implemented, at least in part, in a virtualized cloud architecture, such as an example virtualized cloud architecture **1200** that is illustrated and described herein with reference to FIG. 12. Moreover, aspects of the ICAM system **100** can be implemented, at least in part, through the use of machine learning technologies, such as via an example machine learning system **1300** that is illustrated and described herein with reference to FIG. 13. Those skilled in the will appreciate that the ICAM system **100** can be deployed in various ways on different architectures based upon the needs of a given implementation. Accordingly, the examples set forth herein should not be construed as being limiting to the manner in which the ICAM system **100** is implemented.

(17) In the example illustrated in FIG. 1, the ICAM system **100** can receive an original image **102**. The original image **102** is a digital image. The original image **102** can be a digital photograph, a digital image created by a scanner, a digital image created by software, or other digital image. The original image **102** can depict anything that is capable of classification. As such, the subject matter depicted in the original image **102** is not limited to any particular person, place, or thing. The original image **102** can have any matrix size (e.g., width and height), any pixel size, any resolution (e.g., in terms of pixels per inch “PPI”), any color (e.g., binary, gray-scale, color, or multispectral), any pixel bit depth, and any other image parameters. The original image **102** can be in any file format, some examples of which include, but are not limited to, Tagged Image File Format (“TIFF”), Graphics Interchange Format (“GIF”), Joint Photographic Experts Group (“JPEG”) format, Portable Pix Map (“PPM”), Windows Bitmap (“BMP”), Portable Network Graphics (“PNG”), proprietary file formats, other standardized file formats, and the like.

(18) The illustrated ICAM system **100** includes a plurality of modules, each of which can include instructions that can be executed by one or more processors (see FIGS. 10-12) of the ICAM system **100**. Alternatively, the plurality of modules can be executed by different systems that are operating in communication with one another. In particular, the illustrated ICAM system **100** includes an image resolution reduction module **104**, an ICAM module **106**, an image reconstruction module **108**, an image comparison module **110**, a deep learning image classification (“DLIC”) module **112**, and a classification comparison module **114**. Those skilled in the art will appreciate the numerous ways the disclosed modules can be configured, and as such, the illustrated example described herein should not be construed as being limiting in any way.

(19) The image resolution reduction module **104** can receive the original image **102** and reduce the resolution such that fine details are obscured for analysis. The output of the image resolution reduction module **104** is a reduced resolution image **116**. A pre-determined percentage of resolution reduction or a set resolution target can be used as the basis for reducing the resolution. The reduced resolution image **116** is used so that the ICAM module **106** can perform classification operations faster than deep learning-based classification that is typically used for image classification.

(20) The ICAM module **106** can receive the reduced resolution image **116** from the image resolution reduction module **104** and begin classification operations to generate an ICAM classification **118**. The ICAM classification **118** can be a text-based classification.

(21) In particular, the ICAM module **106** can begin classification of the reduced resolution image **116** by first performing an elimination operation **120** using color as a primary classifier and shape as a secondary classifier. The ICAM module **106** can then perform a slicing operation **122** to slice the reduced resolution image **116** into individual items, and then perform a searching operation **124** to search for common coexisting items associated with the individual items found during the

slicing operation **122** (e.g., ocean waves and a lion normally would not coexist in the same image, but ocean waves and a wooden log would be more likely).

(22) The ICAM module **106** can perform one or more optional classification operations **126**. The optional classification operations **126** can increase the accuracy of the ICAM classification **118** determined by the ICAM module **106**. In some embodiments, the ICAM module **106** can utilize environmental and situational context awareness as one of the optional classification operations **126** to improve classification accuracy. For example, the ICAM module **106** can use the background of the reduced resolution image **116** and its relation to a core subject thereof to determine what is depicted in the original image **102**. The ICAM module **106** can attempt to analyze the cohesiveness of individual elements of the reduced resolution image **116** to better determine the theme of the reduced resolution image **116** and elements that logically go together.

(23) In some embodiments, the ICAM module **106** can utilize textual and/or audible relationships as one of the optional classification operations **126** to improve classification accuracy. For example, the ICAM module **106** can consider any text and/or audio associated with the reduced resolution image **116** with the caveat that this information could be misleading. For example, a clear picture of a tree with text on the picture that identifies the tree as a “flower.” The ICAM module **106** can build a historical trust model for the accuracy of the textual and/or audible description of the images obtained from certain sources.

(24) In some embodiments, the ICAM module **106** can utilize user profile interests as one of the optional classification operations **126**. The classification of an image that depicts an object that is difficult to classify may be aided by a user profile associated with a user who is associated with the image (e.g., in the metadata of the image). In other words, the ICAM module **106** can bias the classification of the reduced resolution image **116** to an object that is associated with an interest of the user. For example, a user profile that indicates boxing as an interest of a user may cause the ICAM module **106** to bias towards boxing-related objects such as boxing gloves.

(25) In some embodiments, the ICAM module **106** can evaluate and determine various objects independently based on relative dimensions and/or mathematical ratios as one of the optional classification operations **126**. The optional classification operations **126** can include other classification operations not explicitly described herein. It is contemplated that, over time, use of the ICAM module **106** may reveal additional optional classification operations **126** that can be used (including experimental use) to improve the accuracy of the ICAM classification **118**.

(26) The image reconstruction module **108** can receive the ICAM classification **118** from the ICAM module **106**. The image reconstruction module **108** can attempt to reconstruct the original image **102** based upon the ICAM classification **118** to create a reconstructed image **128**.

(27) The image comparison module **110** can receive the reconstructed image **128** from the image reconstruction module **108**. The image comparison module **110** can compare the reconstructed image **128** to the original image **102** to determine if the original image was classified accurately. If the image comparison module **110** determines that the comparison is close enough, the image comparison module **110** can determine that the original image **102** was classified accurately. Whether the reconstructed image **128** is close enough to the original image **102** can be determined based upon a similarity threshold. The similarity threshold can be defined as a minimum percentage of matching pixels. For example, if at least 75% of the pixels of the reconstructed image **128** match the original image **102**, then the image comparison module **110** can conclude that the reconstructed image **128** is close enough to the original image **102**. Alternatively, the image comparison module **110** can utilize machine learning to learn correlations among images in terms of coarse details such as shape, subject type (e.g., animal, vehicle, building, person, etc.), and/or other coarse details. For example, two images, one showing a car and the other showing a truck may be considered “close enough,” but two images, one showing a car and the other showing a motorcycle may not be considered “close enough.” The image comparison module **110** can alternatively utilize one or more mathematical formulas such as standard deviation or mean

absolute deviation. Those skilled in the art will appreciate other methods of comparing the reconstructed image **128** and the original image **102**. As such, the aforementioned examples should not be construed as being limiting in any way.

(28) If, however, the image comparison module **110** determines that the comparison is not close enough, the image comparison module **110** can generate and send an adjust resolution instruction **130** to the image resolution reduction module **104**. The adjust resolution instruction **130** can instruct the image resolution reduction module **104** to adjust the resolution of the reduced resolution image **116**. The ICAM module **106**, the image reconstruction module **108**, and the image comparison module **110** can then re-process the reduced resolution image **116**. This process continues until the image comparison module **110** determines that the comparison between the original image **102** and the reconstructed image **128** is close enough.

(29) The DLIC module **112** also processes the original image **102**. In some embodiments, the DLIC module **112** can process the original image **102** in parallel to the ICAM module **106**, although serial processing in which the ICAM module **106** processes the original image **102** before the DLIC module **112**, or vice versa, is also contemplated. The DLIC module **112** can implement a convolutional neural network (“CNN”) **134** to classify the original image **102** and output a DLIC classification **136**. The CNN **134** is an artificial neural network that can be used to analyze and classify the original image **102**. The CNN **134** can use one or more deep learning algorithms to assign weights to various aspects or objects depicted in the original image **102** to differentiate the original image **102** from other images and to assign the DLIC classification **136** to the original image **102**. CNNs are well-known and in common use for image classification tasks. As such, additional details about the CNN **134** are not described herein.

(30) The DLIC module **112** provides the DLIC classification **136** to the classification comparison module **114**. The classification comparison module **114** can also receive the ICAM classification **118**. The classification comparison module **114** can compare the ICAM classification **118** and the DLIC classification **136**. If the ICAM classification **118** and the DLIC classification **136** match, the classification comparison module **114** outputs an ICAM output **138** with the DLIC classification **136**. If, however, the ICAM classification **118** and the DLIC classification **136** do not match, the classification comparison module **114** outputs the ICAM output **138** with the ICAM classification **118**. In some embodiments, the ICAM system **100** can notify one or more other systems and/or devices (not shown) if the DLIC classification **136** is not the same as the ICAM classification **118**, which indicates that the original image **102** contains malicious content (e.g., one or more malicious pixels).

(31) Turning now to FIG. 2, a method **200** for mitigating image classification attacks will be described, according to an illustrative embodiment. It should be understood that the operations of the methods disclosed herein are not necessarily presented in any particular order and that performance of some or all of the operations in an alternative order(s) is possible and is contemplated. The operations have been presented in the demonstrated order for ease of description and illustration. Operations may be added, omitted, and/or performed simultaneously, without departing from the scope of the concepts and technologies disclosed herein.

(32) It also should be understood that the methods disclosed herein can be ended at any time and need not be performed in its entirety. Some or all operations of the methods, and/or substantially equivalent operations, can be performed by execution of computer-readable instructions included on a computer storage media, as defined herein. The term “computer-readable instructions,” and variants thereof, as used herein, is used expansively to include routines, applications, application modules, program modules, programs, components, data structures, algorithms, and the like. Computer-readable instructions can be implemented on various system configurations including single-processor or multiprocessor systems, minicomputers, mainframe computers, personal computers, hand-held computing devices, microprocessor-based, programmable consumer electronics, combinations thereof, and the like.

(33) Thus, it should be appreciated that the logical operations described herein are implemented (1) as a sequence of computer implemented acts or program modules running on a computing system and/or (2) as interconnected machine logic circuits or circuit modules within the computing system. The implementation is a matter of choice dependent on the performance and other requirements of the computing system. Accordingly, the logical operations described herein are referred to variously as states, operations, structural devices, acts, or modules. These states, operations, structural devices, acts, and modules may be implemented in software, in firmware, in special purpose digital logic, and any combination thereof. As used herein, the phrase “cause a processor to perform operations” and variants thereof is used to refer to causing a processor or multiple processors of one or more systems and/or one or more devices disclosed herein to perform one or more operations and/or causing the processor to direct other components of the computing system or device to perform one or more of the operations.

(34) The method **200** begins and proceeds to operation **202**. At operation **202**, the ICAM system **100** executes the ICAM module **106** to classify the original image **102** to determine the ICAM classification **118**. Also at operation **202**, the ICAM module **106** outputs the ICAM classification **118**. Additional details in this regard will be described herein below with reference to FIG. 3.

(35) From operation **202**, the method **200** proceeds to operation **204**. At operation **204**, the ICAM system **100** executes the image reconstruction module **108** in an attempt to reconstruct the original image **102** based on the ICAM classification **118**. The image reconstruction module **108** receives the ICAM classification **118** from the ICAM module **106** and attempts to reconstruct the original image **102** from the ICAM classification **118** to create the reconstructed image **128**. The output of operation **204** is the reconstructed image **128**. At operation **204**, the ICAM system **100** also executes the image comparison module **110** to compare the original image **102** to the reconstructed image **128**. If the original image **102** and the reconstructed image **128** are close enough, the method **200** proceeds to operation **206**. If the original image **102** and the reconstructed image **128** are not close enough, the image comparison module **110** requests, via the adjust resolution instruction **130**, the image resolution reduction module **104** to adjust the resolution of the original image **102** and return to operation **202**. After the image comparison module **110** determines that the original image **102** and the reconstructed image **128** are close enough, the method **200** proceeds to operation **206**. Additional details in this regard will be described herein below with reference to FIG. 4.

(36) At operation **206**, the ICAM system **100** executes the DLIC module **112** to classify the original image **102** to determine the DLIC classification **136**. Also at operation **206**, the DLIC module **112** outputs the DLIC classification **136**. From operation **206**, the method **200** proceeds to operation **208**. At operation **208**, the ICAM system **100** executes the classification comparison module **114** to compare the ICAM classification **118** and the DLIC classification **136**. Also at operation **208**, the ICAM system **100** outputs the appropriate classification based on the comparison. If the DLIC classification **136** matches the ICAM classification **118**, the ICAM system **100** can output the DLIC classification **136**. If the DLIC classification **136** does not match the ICAM classification **118**, the ICAM system **100** can output the ICAM classification **118**. Additional details in this regard will be described herein below with reference to FIG. 5.

(37) From operation **208**, the method **200** proceeds to operation **210**. At operation **210**, the method **200** can end.

(38) Turning now to FIG. 3, a method **300** for classifying the original image **102** to determine the ICAM classification **118** will be described, according to an illustrative embodiment. The method **300** begins and proceeds to operation **302**. At operation **302**, the ICAM system **100** obtains the original image **102**. From operation **302**, the method **300** proceeds to operation **304**. At operation **304**, the ICAM system **100** executes the image resolution reduction module **104** to reduce the resolution of the original image **102** to obscure fine details. The resolution reduction can be based on a pre-determined percentage by which to reduce the resolution of the original image **102**. Alternatively, the resolution reduction can be based on a pre-established target resolution. Other



resolution reduction parameters are contemplated. The image resolution reduction module **104** then outputs the reduced resolution image **116**.

(39) From operation **304**, the method **300** proceeds to operation **306**. At operation **306**, the ICAM system **100** executes the ICAM module **106** to perform the elimination operation **120**. In particular, the ICAM system **100** can perform the elimination operation **120** based on color as a primary classifier and shape as a secondary classifier. From operation **306**, the method **300** proceeds to **308**. At operation **308**, the ICAM system **100** executes the ICAM module **106** to perform the slicing operation **122**. In particular, the ICAM system **100** can slice (i.e., divide) the reduced resolution image **116** into individual items. Also at operation **308**, the ICAM system **100** can search the items for common coexisting items. In some embodiments, the method **300** can then proceed to operation **318** described below. Alternatively, the method **300** can continue by performing one or more of the optional classification operations **126**, which are described below as operations **310**, **312**, **314**, **316**, and **318**. These operations can increase the accuracy of the ICAM classification **118** output by the ICAM module **106**.

(40) From operation **308**, the method **300** proceeds to operation **310**. At operation **310**, the ICAM system **100** executes the ICAM module **106** to perform an environment and situational context awareness check. For example, the ICAM module **106** can use the background of the reduced resolution image **116** and its relation to a core subject thereof to determine what is depicted in the original image **102**. The ICAM module **106** can attempt to analyze the cohesiveness of individual elements of the reduced resolution image **116** to better determine the theme of the reduced resolution image **116** and elements that logically go together.

(41) From operation **310**, the method **300** proceeds to operation **312**. At operation **312**, the ICAM system **100** executes the ICAM module **106** to perform a textual and audible relationship check. In some embodiments, the ICAM module **106** can utilize textual and/or audible relationships as one of the optional classification operations **126** to improve classification accuracy. For example, the ICAM module **106** can consider any text and/or audio associated with the reduced resolution image **116** with the caveat that this information could be misleading. For example, a clear picture of a tree with text on the picture that identifies the tree as a “flower.” The ICAM module **106** can build a historical trust model for the accuracy of the textual and/or audible description of the images obtained from certain sources.

(42) From operation **312**, the method **300** proceeds to operation **314**. At operation **314**, the ICAM system **100** executes the ICAM module **106** to perform user profile biasing. The classification of an image that depicts an object that is difficult to classify may be aided by a user profile associated with a user who is associated with the image (e.g., in the metadata of the image). In other words, the ICAM module **106** can bias the classification of the reduced resolution image **116** to an object that is associated with an interest of the user. For example, a user profile that indicates boxing as an interest of a user may cause the ICAM module **106** to bias towards boxing-related objects such as boxing gloves.

(43) From operation **314**, the method **300** proceeds to operation **316**. At operation **316**, the ICAM system executes the ICAM module to perform a relative dimension and mathematical ratio analysis. In some embodiments, the ICAM module **106** can evaluate and determine various objects independently based on relative dimensions and/or mathematical ratios as one of the optional classification operations **126**. The optional classification operations **126** can include other classification operations not explicitly described herein. It is contemplated that, over time, use of the ICAM module **106** may reveal additional optional classification operations **126** that can be used (including experimental use) to improve the accuracy of the ICAM classification **118**.

(44) From operation **316**, the method **300** proceeds to operation **318**. At operation **318**, the ICAM system **100** executes the ICAM module **106** to classify the original image **102** and provide a textual output of the ICAM classification **118**.

(45) From operation **318**, the method **300** proceeds to operation **320**. The method **300** can end at

operation 320.

(46) Turning now to FIG. 4, a method 400 for reconstructing the original image 102 from the ICAM classification 118 will be described, according to an illustrative embodiment. The method 400 begins and proceeds to operation 402. At operation 402, the ICAM system 100 executes the image reconstruction module 108 to obtain the ICAM classification 118. From operation 402, the method 400 proceeds to operation 404. At operation 404, the ICAM system 100 executes the image reconstruction module 108 to create the reconstructed image 128 based on the ICAM classification 118.

(47) From operation 404, the method 400 proceeds to operation 406. At operation 406, the ICAM system 100 executes the image comparison module 110 to compare the original image 102 to the reconstructed image 128. From operation 406, the method 400 proceeds to operation 408. At operation 408, the ICAM system 100 executes the image comparison module 110 to determine if the reconstructed image 128 is close enough to the original image 102. If the image comparison module 110 determines that the reconstructed image 128 is close enough to the original image 102, the method 400 can proceed to operation 410. At operation 410, the image comparison module 110 determines that the original image 102 was processed accurately. From operation 410, the method 400 proceeds to operation 412. The method 400 can end at operation 412.

(48) Returning to operation 408, if the image comparison module 110 determines that the reconstructed image 128 is not close enough to the original image 102, the method 400 proceeds to operation 414. At operation 414, the ICAM system 100 executes the image resolution reduction module 104 to adjust the resolution of the original image 102. From operation 414, the method 400 returns to operation 306 of the method 300 shown in FIG. 3, which is described above.

(49) Turning now to FIG. 5, a method 500 for performing image processing to classify the original image 102 using the DLIC module 112 and compare the resultant DLIC classification 136 to the ICAM classification 118 will be described, according to an illustrative embodiment. The method 500 begins and proceeds to operation 502. At operation 502, the ICAM system 100 executes the DLIC module 112 to perform image classification via the CNN 134 to determine the DLIC classification 136 of the original image 102.

(50) From operation 502, the method 500 proceeds to operation 504. At operation 504, the ICAM system 100 executes the classification comparison module 114 to compare the ICAM classification 118 to the DLIC classification 136. From operation 504, the method 500 proceeds to operation 506. At operation 506, the ICAM system 100 executes the classification comparison module 114 to determine if the ICAM classification 118 and the DLIC classification 136 match. If, at operation 506, the classification comparison module 114 determines that the ICAM classification 118 and the DLIC classification 136 match, the method 500 proceeds to operation 508. At operation 508, the classification comparison module 114 provide a textual output of the DLIC classification 136. From operation 508, the method 500 proceeds to operation 510. The method 500 can end at operation 510.

(51) Returning to operation 506, if the classification comparison module 114 determines that the ICAM classification 118 and the DLIC classification 136 do not match, the method 500 proceeds to operation 512. At operation 512, the classification comparison module presents the ICAM classification. From operation 512, the method 500 proceeds to operation 514. At operation 514, ICAM system 100 can perform a remedial action. For example, the ICAM system 100 can notify a user, owner, or other entity associated with the ICAM system 100 that the DLIC module 112 has been compromised. From operation 514, the method 500 can proceed to operation 510. The method 500 can end at operation 510.

(52) Turning now to FIG. 6, an audio attack mitigation (“AAM”) system 600 will be described, according to an illustrative embodiment. The AAM system 600 can be implemented, at least in part, in a computer system, such as an example computer system 1000 that is illustrated and described with reference to FIG. 10. The AAM system 600 alternatively can be implemented, at

least in part, in a containerized architecture, such as an example containerized cloud architecture **1100** that is illustrated and described herein with reference to FIG. **11**. The AAM system **600** can be implemented, at least in part, in a virtualized cloud architecture, such as an example virtualized cloud architecture **1200** that is illustrated and described herein with reference to FIG. **12**. Moreover, aspects of the AAM system **600** can be implemented, at least in part, through the use of machine learning technologies, such as via an example machine learning system **1300** that is illustrated and described herein with reference to FIG. **13**. Those skilled in the will appreciate that the AAM system **600** can be deployed in various ways on different architectures based upon the needs of a given implementation. Accordingly, the examples set forth herein should not be construed as being limiting to the manner in which the AAM system **600** is implemented.

(53) In the example illustrated in FIG. **6**, the AAM system **600** can receive an original digital audio signal (“original audio”) **602**. The original audio **602** can be in any file format, some examples of which include, but are not limited to, pulse-code modulation (“PCM”), Waveform Audio File Format (“WAV”), Audio Interchange File Format (“AIFF”), Moving Pictures Expert Group (“MPEG”) Audio Layer 3 (“MP3”), MPEG Audio Layer 4 (“MP4”), Advanced Audio Coding (“AAC”), Windows Media Audio (“WMA”), Free Lossless Audio Codec (“FLAC”), Apple Lossless Audio Codec (“ALAC”), proprietary file formats, other standardized file formats, and the like.

(54) The original audio **602** can be compromised. An attacker may inject a malicious undetectable waveform into the original audio **602** such that a receiver will decode and transcribe words that did not exist in the original audio **602**. The original audio **602** may be used by a destination system **604** to perform an action. For example, the destination system **604** might be an autonomous vehicle or system thereof. In this example, an audio command such as “stop vehicle” may be compromised with a malicious undetectable waveform that causes the audio command to be transcribed instead as “accelerate vehicle.” This may result in the vehicle crashing and injuring or killing the passenger(s). As another example, the destination system **604** might be a voice-enabled home assistant that enables a user to control smart home devices such as a smart lock. In this example, an audio command such as “lock front door” may be compromised with a malicious undetectable waveform that causes the audio command to be transcribed instead as “unlock front door.” This may expose the user's home to a robbery or other crime.

(55) The AAM system **600** can receive the original audio **602** and provide the original audio **602** to a multi-rate sampler and text generator module **606**. The multi-rate sampler and text generator module **606** can sample the original audio **602** at multiple bit depths (e.g., 8-bit, 16-bit, 24-bit, etc.) and/or sampling rates (e.g., 44.1 kHz, 48 kHz, 96 kHz, 192 kHz, etc.) to create multiple audio samples of the original audio **602**. The multi-rate sampler and text generator module **606** transcribes the audio samples into text samples **608A-608N**.

(56) A text comparison module.sub.1 **610A** can receive and compare the text samples **608A-608N** to determine if the text samples **608A-608N** match. If the text comparison module.sub.1 **610A** determines that the text samples **608A-608N** do not match (i.e., there is some discrepancy among the text samples **608A-608N**), the text comparison module.sub.1 **610A** can generate an alarm.sub.1 **612A** directed to the destination system **604**. If the text comparison module.sub.1 **610A** determines that the text samples **608A-608N** do match, the text comparison module.sub.1 **610A** can generate a text outputs **614A** directed to a final text comparison module **616**.

(57) The AAM system **600** can also process the original audio **602** via a multi-level audio enhancer and text generator module **618** and a multi-level audio compressor and text generator module **620**. This process can be performed in parallel to the processing described above. Alternatively, the processes can be serialized. The multi-level audio enhancer and text generator module **618** can receive the original audio **602** and enhance the original audio **602** to increase clarity. For example, the multi-level audio enhancer and text generator module **618** can increase the clarity of voice audio by filling in any gaps in the waveform, reconstructing any distortion present in the

waveform, and smoothing out the waveform. The multi-level audio enhancer and text generator module **618** can then transcribe the enhanced audio into enhanced audio text **622**. The multi-level audio compressor and text generator module **620** can receive the original audio **602** and compress the original audio **602** into lossy compressed audio. The compression process can utilize any lossy compression algorithm to reduce the file size of the original audio **602** by omitting fine details in the audio waveform. The multi-level audio compressor and text generator module **620** can then transcribe the compressed audio into compressed audio text **624**.

(58) A text comparison module.sub.2 **610B** can receive and compare the enhanced audio text **622** and the compressed audio text **624** to determine if they match. The illustrated embodiment shows two separate text comparison modules **610A**, **610B**, although a single text comparison module **610** is contemplated. If the text comparison module.sub.2 **610B** determines that the enhanced audio text **622** and the compressed audio text **624** do not match (i.e., there is some discrepancy between the enhanced audio text **622** and the compressed audio text **624**), the text comparison module.sub.2 **610B** can generate an alarm.sub.2 **612B** directed to the destination system **604**. If the text comparison module.sub.2 **610B** determines that the enhanced audio text **622** and the compressed audio text **624** do match, the text comparison module.sub.2 **610B** can generate a text output.sub.2 **614B** directed to the final text comparison module **616**.

(59) The final text comparison module **616** can compare the text output.sub.1 **614A** and the text output.sub.2 **614B**. If the final text comparison module **616** determines that the text output.sub.1 **614A** and the text output.sub.2 **614B** match, the final text comparison module **616** can determine that the original audio **602** has not been compromised and can generate a final text output **626** directed to the destination system **604**. If the final text comparison module **616** determines that the text output.sub.1 **614A** and the text output.sub.2 **614B** do not match, the final text comparison module **616** can determine that the original audio **602** has been compromised and can generate an alarm.sub.3 **612C** directed to the destination system **604**.

(60) Turning now to FIG. 7, a method **700** for comparing the text samples **608A**, **608B** generated from multiple audio samples will be described, according to an illustrative embodiment. The method **700** begins and proceeds to operation **702**. At operation **702**, the multi-rate sampler and text generator module **606** receives the original audio **602**. From operation **702**, the method **700** proceeds to operation **704**. At operation **704**, the multi-rate sampler and text generator module **606** samples the original audio **602** at multiple bit depths (e.g., 8-bit, 16-bit, 24-bit, etc.) and/or sampling rates (e.g., 44.1 kHz, 48 kHz, 96 kHz, 192 kHz, etc.) to create multiple audio samples of the original audio **602**. Also at operation **704**, the multi-rate sampler and text generator module **606** generates and outputs the text samples **608A-608N** generated from the audio samples.

(61) From operation **704**, the method **700** proceeds to operation **706**. At operation **706**, the text comparison module.sub.1 **610A** compares the text samples **608A-608N**. From operation **706**, the method **700** proceeds to operation **708**. At operation **708**, the text comparison module.sub.1 **610A** determines if the text samples **608A-608N** match. If the text comparison module.sub.1 **610A** determines that the text samples **608A-608N** match, the method **700** proceeds to operation **710**. At operation **710**, the text comparison module.sub.1 **610A** generates the text output.sub.1 **614A** and provides the text output.sub.1 **614A** to the final text comparison module **616**. From operation **710**, the method **700** proceeds to operation **902** of the method **900** described herein below. If the text comparison module.sub.1 **610A** determines that the text samples **608A-608N** do not match (i.e., there is some discrepancy among the texts), the method **700** proceeds to operation **712**. At operation **712**, the text comparison module.sub.1 **610A** generates and provides the alarm.sub.1 **612A** to the destination system **604**. From operation **712**, the method **700** proceeds to operation **714**. At operation **714**, the method **700** can end.

(62) Turning now to FIG. 8, a method **800** for comparing text generated from the enhanced audio text **622** and the compressed audio text **624** will be described, according to an illustrative embodiment. The method **800** begins and proceeds to operation **802**. At operation **802**, both the

multi-level audio enhancer and text generator module **618** and the multi-level audio compressor and text generator module **620** receive the original audio **602**. From operation **802**, the method **800** proceeds to operation **804**. At operation **804**, the multi-level audio enhancer and text generator module **618** performs audio enhancement on the original audio **602** to create enhanced audio. Also at operation **804**, the multi-level audio enhancer and text generator module **618** generates and outputs the enhanced audio text **622**.

(63) From operation **804**, the method **800** proceeds to operation **806**. At operation **806**, the multi-level audio compressor and text generator module **620** performs audio compression on the original audio **602**. Also at operation **806**, the multi-level audio compression and text generator module **620** generates and outputs the compressed audio text **624**.

(64) From operation **806**, the method proceeds to operation **808**. At operation **808**, the text comparison module.sub.2 **610B** compares the enhanced audio text **622** and the compressed audio text **624**. From operation **808**, the method **800** proceeds to operation **810**. At operation **810**, the text comparison module.sub.2 **610B** determines if the enhanced audio text **622** and the compressed audio text match **624** match. If the text comparison module.sub.2 **610B** determines that the enhanced audio text **622** and the compressed audio text **624** match, the method **800** proceeds to operation **812**. At operation **812**, the text comparison module.sub.2 **610B** generates the text output.sub.2 **614B** and provides the text output.sub.2 **614B** to the final text comparison module **616**. From operation **812**, the method **800** proceeds to operation **902** of the method **900** described herein below. If the text comparison module.sub.2 **610B** determines that the enhanced audio text **622** and the compressed audio text **624** do not match (i.e., there is some discrepancy among the texts), the method **800** proceeds to operation **814**. At operation **814**, the text comparison module.sub.2 **610B** generates and provides the alarm.sub.2 **612B** to the destination system **604**. From operation **814**, the method **800** proceeds to operation **816**. The method **800** can end at operation **816**.

(65) Turning now to FIG. **9**, a method **900** for comparing the text outputs **614A**, **614B** obtained from the methods **700**, **800** will be described, according to an illustrative embodiment. The method **900** begins as a continuation from operation **710** (see FIG. **7**) and operation **812** (see FIG. **8**) described above. At operation **902**, the final text comparison module **616** receives and compares the text outputs **614A**, **614B**. From operation **902**, the method **900** proceeds to operation **904**. At operation **904**, the final text comparison module **616** determines if the text outputs **614A**, **614B** match. If the final text comparison module **616** determines that the text outputs **614A**, **614B** match, then the method **900** proceeds to operation **906**. At operation **906**, the final text comparison module **616** generates the final text output **626** and provides the final text output **626** to the destination system **604**. From operation **906**, the method **900** proceeds to operation **908**. The method **900** can end at operation **908**.

(66) Returning to operation **904**, if the final text comparison module **616** determines that the text outputs **614A**, **614B** do not match, then the method **900** proceeds to operation **910**. At operation **910**, the final text comparison module **616** generates and presents the alarm.sub.3 **612C** to the destination system **604**. From operation **910**, the method **900** proceeds to operation **908**. The method **900** can end at operation **908**.

(67) Turning now to FIG. **10**, a block diagram illustrating a computer system **1000** configured to provide the functionality described herein in accordance with various embodiments of the concepts and technologies disclosed herein. In some embodiments, the ICAM system **100** and/or the AAM system **600** can be configured the same as or similar to the computer system **1000**. The computer system **1000** includes a processing unit **1002**, a memory **1004**, one or more user interface devices **1006**, one or more input/output (“I/O”) devices **1008**, and one or more network devices **1010**, each of which is operatively connected to a system bus **1012**. The bus **1012** enables bi-directional communication between the processing unit **1002**, the memory **1004**, the user interface devices **1006**, the I/O devices **1008**, and the network devices **1010**.

(68) The processing unit **1002** may be a standard central processor that performs arithmetic and logical operations, a more specific purpose programmable logic controller (“PLC”), a programmable gate array, or other type of processor known to those skilled in the art and suitable for controlling the operation of the server computer. The processing unit **1002** can be a single processing unit or a multiple processing unit that includes more than one processing component. Processing units are generally known, and therefore are not described in further detail herein.

(69) The memory **1004** communicates with the processing unit **1002** via the system bus **1012**. The memory **1004** can include a single memory component or multiple memory components. In some embodiments, the memory **1004** is operatively connected to a memory controller (not shown) that enables communication with the processing unit **1002** via the system bus **1012**. The memory **1004** includes an operating system **1014** and one or more program modules **1016**. The operating system **1014** can include, but is not limited to, members of the WINDOWS, WINDOWS CE, and/or WINDOWS MOBILE families of operating systems from MICROSOFT CORPORATION, the LINUX family of operating systems, the SYMBIAN family of operating systems from SYMBIAN LIMITED, the BREW family of operating systems from QUALCOMM CORPORATION, the MAC OS, iOS, and/or LEOPARD families of operating systems from APPLE CORPORATION, the FREEBSD family of operating systems, the SOLARIS family of operating systems from ORACLE CORPORATION, other operating systems, and the like.

(70) The program modules **1016** may include various software and/or program modules described herein. In some embodiments, the program modules **1016** in the ICAM system **100** configured like the computer system **1000** can include, for example, the image resolution reduction module **104**, the ICAM module **106**, the image reconstruction module **108**, the image comparison module **110**, the DLIC module **112**, a classification comparison module **114**, or a combination thereof. In some embodiments, the program modules **1016** in the AAM system **600** configured like the computer system **1000** can include, for example, the multi-rate sampler and text generator module **606**, the text comparison modules **610A**, **610B**, the final text comparison module **616**, or a combination thereof. In some embodiments, multiple implementations of the computer system **1000** can be used, wherein each implementation is configured to execute one or more of the program modules **1016**. The program modules **1016** and/or other programs can be embodied in computer-readable media containing instructions that, when executed by the processing unit **1002**, perform the methods described herein. According to embodiments, the program modules **1016** may be embodied in hardware, software, firmware, or any combination thereof. Although not shown in FIG. **10**, it should be understood that the memory **1004** also can be configured to store the original image **102**, the reduced resolution image **116**, the ICAM classification **118**, the reconstructed image **128**, the adjust resolution instruction **130**, the DLIC classification **136**, the CNN **134**, the ICAM output **138**, combinations thereof, and/or other data disclosed herein.

(71) By way of example, and not limitation, computer-readable media may include any available computer storage media or communication media that can be accessed by the computer system **1000**. Communication media includes computer-readable instructions, data structures, program modules, or other data in a modulated data signal such as a carrier wave or other transport mechanism and includes any delivery media. The term “modulated data signal” means a signal that has one or more of its characteristics changed or set in a manner as to encode information in the signal. By way of example, and not limitation, communication media includes wired media such as a wired network or direct-wired connection, and wireless media such as acoustic, RF, infrared and other wireless media. Combinations of any of the above should also be included within the scope of computer-readable media.

(72) Computer storage media includes volatile and non-volatile, removable and non-removable media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, or other data. Computer storage media includes, but is not limited to, RAM, ROM, Erasable Programmable ROM (“EPROM”),

Electrically Erasable Programmable ROM (“EEPROM”), flash memory or other solid state memory technology, CD-ROM, digital versatile disks (“DVD”), or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store the desired information and which can be accessed by the computer system **1000**. In the claims, the phrase “computer storage medium,” “computer-readable storage medium,” and variations thereof does not include waves or signals per se and/or communication media, and therefore should be construed as being directed to “non-transitory” media only.

(73) The user interface devices **1006** may include one or more devices with which a user accesses the computer system **1000**. The user interface devices **1006** may include, but are not limited to, computers, servers, personal digital assistants, cellular phones, or any suitable computing devices. The I/O devices **1008** enable a user to interface with the program modules **1016**. In one embodiment, the I/O devices **1008** are operatively connected to an I/O controller (not shown) that enables communication with the processing unit **1002** via the system bus **1012**. The I/O devices **1008** may include one or more input devices, such as, but not limited to, a keyboard, a mouse, or an electronic stylus. Further, the I/O devices **1008** may include one or more output devices, such as, but not limited to, a display screen or a printer.

(74) The network devices **1010** enable the computer system **1000** to communicate with other networks or remote systems via a network **1018**. Examples of the network devices **1010** include, but are not limited to, a modem, a radio frequency (“RF”) or infrared (“IR”) transceiver, a telephonic interface, a bridge, a router, or a network card. The network **1018** may include a wireless network such as, but not limited to, a Wireless Local Area Network (“WLAN”) such as a WI-FI network, a Wireless Wide Area Network (“WWAN”), a Wireless Personal Area Network (“WPAN”) such as BLUETOOTH, a Wireless Metropolitan Area Network (“WMAN”) such as WiMAX network, or a cellular network. Alternatively, the network **1018** may be a wired network such as, but not limited to, a Wide Area Network (“WAN”) such as the Internet, a Local Area Network (“LAN”) such as the Ethernet, a wired Personal Area Network (“PAN”), or a wired Metropolitan Area Network (“MAN”).

(75) Turning now to FIG. **11**, a block diagram illustrating an exemplary containerized cloud architecture **1100** capable of implementing, at least in part, aspects of the concepts and technologies disclosed herein will be described, according to an illustrative embodiment. In some embodiments, the ICAM system **100** and/or the audio attack mitigation system **600**, at least in part, is implemented in the containerized cloud architecture **1100**. The illustrated containerized cloud architecture **1100** includes a first host (“host.sub.1”) **1102A** and a second host (“host.sub.2”) **1102B** (at times referred to herein collectively as hosts **1102** or individually as host **1102**) that can communicate via an overlay network **1104**. Although two hosts **1102** are shown, the containerized cloud architecture **1100** can support any number of hosts **1102**. The overlay network **1104** can enable communication among hosts **1102** in the same cloud network or hosts **1102** across different cloud networks. Moreover, the overlay network **1104** can enable communication among hosts **1102** owned and/or operated by the same or different entities.

(76) The illustrated host.sub.1 **1102A** includes a host hardware **1106A**, a host operating system.sub.1 **1108A**, a DOCKER engine.sub.1 **1110A**, a bridge network.sub.1 **1112A**, container.sub.A-1 through container.sub.N-1 **1114A1-1114N1**, and microservice.sub.A-1 through microservice.sub.N-1 **1116A1-1116N1**. Similarly, the illustrated host.sub.2 **1102B** includes a host hardware.sub.2 **1106B**, a host operating system.sub.2 **1108B**, a DOCKER engine.sub.2 **1110B**, a bridge network.sub.2 **1112B**, container.sub.A-2 through container.sub.N-2 **1114A2-1114N2**, and microservice.sub.A-2 through microservice.sub.N-2 **1116A2-1116N2**.

(77) The host hardware **1106A** and the host hardware **1106B** (at times referred to herein collectively or individually as host hardware **1106**) can be implemented as bare metal hardware such as one or more physical servers. The host hardware **1106** alternatively can be implemented

using hardware virtualization. In some embodiments, the host hardware **1106** can include compute resources, memory resources, and other hardware resources. These resources can be virtualized according to known virtualization techniques. A virtualization cloud architecture **1200** is described herein with reference to FIG. **12**. Although the containerized cloud architecture **1100** and the virtualization cloud architecture **1200** are described separately, these architectures can be combined to provide a hybrid containerized/virtualized cloud architecture. Those skilled in the art will appreciate that the disclosed cloud architectures are simplified for ease of explanation and can be altered as needed for any given implementation without departing from the scope of the concepts and technologies disclosed herein. As such, the containerized cloud architecture **1100** and the virtualized cloud architecture **1200** should not be construed as being limiting in any way.

(78) Compute resources can include one or more hardware components that perform computations to process data and/or to execute computer-executable instructions. For example, the compute resources can execute instructions of the host operating system **1108A** and the host operating systems **1108B** (at times referred to herein collectively as host operating systems **1108** or individually as host operating system **1108**), the containers **1114A1-1114N1** and the containers **1114A2-1114N2** (at times referred to herein collectively as containers **1114** or individually as container **1114**), and the microservices **1116A1-1116N1** and the microservices **1116A1-1116N1** (at times referred to herein collectively as microservices **1116** or individually as microservice **1116**).

(79) The compute resources of the host hardware **1106** can include one or more central processing units (“CPUs”) configured with one or more processing cores. The compute resources can include one or more graphics processing unit (“GPU”) configured to accelerate operations performed by one or more CPUs, and/or to perform computations to process data, and/or to execute computer-executable instructions of one or more application programs, operating systems, and/or other software that may or may not include instructions particular to graphics computations. In some embodiments, the compute resources can include one or more discrete GPUs. In some other embodiments, the compute resources can include CPU and GPU components that are configured in accordance with a co-processing CPU/GPU computing model, wherein the sequential part of an application executes on the CPU and the computationally-intensive part is accelerated by the GPU. The compute resources can include one or more system-on-chip (“SoC”) components along with one or more other components, including, for example, one or more memory resources, and/or one or more other resources. In some embodiments, the compute resources can be or can include one or more SNAPDRAGON SoCs, available from QUALCOMM; one or more TEGRA SoCs, available from NVIDIA; one or more HUMMINGBIRD SoCs, available from SAMSUNG; one or more Open Multimedia Application Platform (“OMAP”) SoCs, available from TEXAS INSTRUMENTS; one or more customized versions of any of the above SoCs; and/or one or more proprietary SoCs. The compute resources can be or can include one or more hardware components architected in accordance with an advanced reduced instruction set computing (“RISC”) (“ARM”) architecture, available for license from ARM HOLDINGS. Alternatively, the compute resources can be or can include one or more hardware components architected in accordance with an x86 architecture, such as an architecture available from INTEL CORPORATION, and others. Those skilled in the art will appreciate the implementation of the compute resources can utilize various computation architectures, and as such, the compute resources should not be construed as being limited to any particular computation architecture or combination of computation architectures, including those explicitly disclosed herein.

(80) The memory resources of the host hardware **1106** can include one or more hardware components that perform storage operations, including temporary or permanent storage operations. In some embodiments, the memory resource(s) include volatile and/or non-volatile memory implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, or other data disclosed herein. Computer storage media includes, but is not limited to, random access memory (“RAM”), read-only memory



(“ROM”), Erasable Programmable ROM (“EPROM”), Electrically Erasable Programmable ROM (“EEPROM”), flash memory or other solid state memory technology, CD-ROM, digital versatile disks (“DVD”), or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store data and which can be accessed by the compute resources.

(81) The other resource(s) of the host hardware **1106** can include any other hardware resources that can be utilized by the compute resources(s) and/or the memory resource(s) to perform operations described herein. The other resource(s) can include one or more input and/or output processors (e.g., network interface controller or wireless radio), one or more modems, one or more codec chipset, one or more pipeline processors, one or more fast Fourier transform (“FFT”) processors, one or more digital signal processors (“DSPs”), one or more speech synthesizers, and/or the like.

(82) The host operating systems **1108** can be proprietary, open source, or closed source. In some embodiments, the host operating systems **1108** can be or can include one or more container operating systems designed specifically to host containers such as the containers **1114**. For example, the host operating systems **1108** can be or can include FEDORA COREOS (available from RED HAT, INC), RANCHEROS (available from RANCHER), and/or BOTTLEROCKET (available from Amazon Web Services). In some embodiments, the host operating systems **1108** can be or can include one or more members of the WINDOWS family of operating systems from MICROSOFT CORPORATION (e.g., WINDOWS SERVER), the LINUX family of operating systems (e.g., CENTOS, DEBIAN, FEDORA, ORACLE LINUX, RHEL, SUSE, and UBUNTU), the SOLARIS family of operating systems from ORACLE CORPORATION, other operating systems, and the like.

(83) The containerized cloud architecture **1100** can be implemented utilizing any containerization technologies. Presently, open-source container technologies, such as those available from DOCKER, INC., are the most widely used, and it appears will continue to be for the foreseeable future. For this reason, the containerized cloud architecture **1100** is described herein using DOCKER container technologies available from DOCKER, INC., such as the DOCKER engines **1110**. Those skilled in the art will appreciate that other container technologies, such as KUBERNETES may also be applicable to implementing the concepts and technologies disclosed herein, and as such, the containerized cloud architecture **1100** is not limited to DOCKER container technologies. Moreover, although open-source container technologies are most widely used, the concepts and technologies disclosed here may be implemented using proprietary technologies or closed source technologies.

(84) The DOCKER engines **1110** are based on open source containerization technologies available from DOCKER, INC. The DOCKER engines **1110** enable users (not shown) to build and containerize applications. The full breadth of functionality provided by the DOCKER engines **1110** and associated components in the DOCKER architecture are beyond the scope of the present disclosure. As such, the primary functions of the DOCKER engines **1110** will be described herein in brief, but this description should not be construed as limiting the functionality of the DOCKER engines **1110** or any part of the associated DOCKER architecture. Instead, those skilled in the art will understand the implementation of the DOCKER engines **1110** and other components of the DOCKER architecture to facilitate building and containerizing applications within the containerized cloud architecture **1100**.

(85) The DOCKER engine **1110** functions as a client-server application executed by the host operating system **1108**. The DOCKER engine **1110** provides a server with a daemon process along with application programming interfaces (“APIs”) that specify interfaces that applications can use to communicate with and instruct the daemon to perform operations. The DOCKER engine **1110** also provides a command line interface (“CLI”) that uses the APIs to control and interact with the daemon through scripting and/or CLI commands. The daemon can create and manage objects such as images, containers, networks, and volumes. Although a single DOCKER engine **1110** is

illustrated in each of the hosts **1102**, multiple DOCKER engines **1110** are contemplated. The DOCKER engine(s) **1110** can be run in swarm mode.

(86) The bridge networks **1112** enable the containers **1114** connected to the same bridge network to communicate. For example, the bridge networks **1112A** enables communication among the containers **1114A1-1114N1**, and the bridge network.sub.2 **1112B** enables communication among the containers **1114A2-1114N2**. In some embodiments, the bridge networks **1112** are software network bridges implemented via the DOCKER bridge driver. The DOCKER bridge driver enables default and user-defined network bridges. The containers **1114** are runtime instances of images. The containers **1114** are described herein specifically as DOCKER containers, although other containerization technologies are contemplated as noted above. Each container **1114** can include an image, an execution environment, and a standard set of instructions

(87) The microservices **1116** are applications that provide a single function. In some embodiments, each of the microservices **1116** is provided by one of the containers **1114**, although each of the containers **1114** may contain multiple microservices **1116**. For example, the microservices **1116** can include, but are not limited, to server, database, and other executable applications to be run in an execution environment provided by a container **1114**. The microservices **1116** can provide any type of functionality, and therefore all the possible functions cannot be listed herein. Those skilled in the art will appreciate the use of the microservices **1116** along with the containers **1114** to improve many aspects of the containerized cloud architecture **1100**, such as reliability, security, agility, and efficiency, for example.

(88) Turning now to FIG. **12**, a block diagram illustrating an example virtualized cloud architecture **1200** and components thereof will be described, according to an exemplary embodiment. The virtualized cloud architecture **1200** can be utilized to implement various elements disclosed herein. In some embodiments, the ICAM system **100** and/or the AAM system **600**, at least in part, is implemented in the virtualized cloud architecture **1200**.

(89) The virtualized cloud architecture **1200** is a shared infrastructure that can support multiple services and network applications. The illustrated virtualized cloud architecture **1200** includes a hardware resource layer **1202**, a control layer **1204**, a virtual resource layer **1206**, and an application layer **1208** that work together to perform operations as will be described in detail herein.

(90) The hardware resource layer **1202** provides hardware resources, which, in the illustrated embodiment, include one or more compute resources **1210**, one or more memory resources **1212**, and one or more other resources **1214**. The compute resource(s) **1210** can include one or more hardware components that perform computations to process data, and/or to execute computer-executable instructions of one or more application programs, operating systems, and/or other software. The compute resources **1210** can include one or more central processing units (“CPUs”) configured with one or more processing cores. The compute resources **1210** can include one or more graphics processing unit (“GPU”) configured to accelerate operations performed by one or more CPUs, and/or to perform computations to process data, and/or to execute computer-executable instructions of one or more application programs, operating systems, and/or other software that may or may not include instructions particular to graphics computations. In some embodiments, the compute resources **1210** can include one or more discrete GPUs. In some other embodiments, the compute resources **1210** can include CPU and GPU components that are configured in accordance with a co-processing CPU/GPU computing model, wherein the sequential part of an application executes on the CPU and the computationally-intensive part is accelerated by the GPU. The compute resources **1210** can include one or more system-on-chip (“SoC”) components along with one or more other components, including, for example, one or more of the memory resources **1212**, and/or one or more of the other resources **1214**. In some embodiments, the compute resources **1210** can be or can include one or more SNAPDRAGON SoCs, available from QUALCOMM; one or more TEGRA SoCs, available from NVIDIA; one or more

HUMMINGBIRD SoCs, available from SAMSUNG; one or more Open Multimedia Application Platform (“OMAP”) SoCs, available from TEXAS INSTRUMENTS; one or more customized versions of any of the above SoCs; and/or one or more proprietary SoCs. The compute resources **1210** can be or can include one or more hardware components architected in accordance with an advanced reduced instruction set computing (“RISC”) machine (“ARM”) architecture, available for license from ARM HOLDINGS. Alternatively, the compute resources **1210** can be or can include one or more hardware components architected in accordance with an x86 architecture, such an architecture available from INTEL CORPORATION of Mountain View, California, and others. Those skilled in the art will appreciate the implementation of the compute resources **1210** can utilize various computation architectures, and as such, the compute resources **1210** should not be construed as being limited to any particular computation architecture or combination of computation architectures, including those explicitly disclosed herein.

(91) The memory resource(s) **1212** can include one or more hardware components that perform storage operations, including temporary or permanent storage operations. In some embodiments, the memory resource(s) **1212** include volatile and/or non-volatile memory implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, or other data disclosed herein. Computer storage media includes, but is not limited to, random access memory (“RAM”), read-only memory (“ROM”), Erasable Programmable ROM (“EPROM”), Electrically Erasable Programmable ROM (“EEPROM”), flash memory or other solid state memory technology, CD-ROM, digital versatile disks (“DVD”), or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store data and which can be accessed by the compute resources **1210**.

(92) The other resource(s) **1214** can include any other hardware resources that can be utilized by the compute resources(s) **1210** and/or the memory resource(s) **1212** to perform operations described herein. The other resource(s) **1214** can include one or more input and/or output processors (e.g., network interface controller or wireless radio), one or more modems, one or more codec chipset, one or more pipeline processors, one or more fast Fourier transform (“FFT”) processors, one or more digital signal processors (“DSPs”), one or more speech synthesizers, and/or the like.

(93) The hardware resources operating within the hardware resources layer **1202** can be virtualized by one or more virtual machine monitors (“VMMs”) **1216A-1216N** (also known as “hypervisors;” hereinafter “VMMs **1216**”) operating within the control layer **1204** to manage one or more virtual resources that reside in the virtual resource layer **1206**. The VMMs **1216** can be or can include software, firmware, and/or hardware that alone or in combination with other software, firmware, and/or hardware, manages one or more virtual resources operating within the virtual resource layer **1206**.

(94) The virtual resources operating within the virtual resource layer **1206** can include abstractions of at least a portion of the compute resources **1210**, the memory resources **1212**, the other resources **1214**, or any combination thereof. These abstractions are referred to herein as virtual machines (“VMs”). In the illustrated embodiment, the virtual resource layer **1206** includes VMs **1218A-1218N** (hereinafter “VMs **1218**”). Each of the VMs **1218** can execute one or more applications **1220A-1220N** in the application layer **1208**.

(95) Turning now to FIG. **13**, a machine learning system **1300** capable of implementing aspects of the embodiments disclosed herein will be described. In some embodiments, aspects of the ICAM system **100** and/or the AAM system **600** can use of machine learning and/or artificial intelligence applications. Accordingly, the ICAM system **100** and/or the AAM system **600** can include the machine learning system **1300** or can be in communication with the machine learning system **1300**.

(96) The illustrated machine learning system **1300** includes one or more machine learning models **1302**. The machine learning models **1302** can include supervised and/or semi-supervised learning

models. The machine learning model(s) **1302** can be created by the machine learning system **1300** based upon one or more machine learning algorithms **1304**. The machine learning algorithm(s) **1304** can be any existing, well-known algorithm, any proprietary algorithms, or any future machine learning algorithm. Some example machine learning algorithms **1304** include, but are not limited to, neural networks, gradient descent, linear regression, logistic regression, linear discriminant analysis, classification tree, regression tree, Naive Bayes, K-nearest neighbor, learning vector quantization, support vector machines, and the like. Classification and regression algorithms might find particular applicability to the concepts and technologies disclosed herein. Those skilled in the art will appreciate the applicability of various machine learning algorithms **1304** based upon the problem(s) to be solved by machine learning via the machine learning system **1300**.

(97) The machine learning system **1300** can control the creation of the machine learning models **1302** via one or more training parameters. In some embodiments, the training parameters are selected modelers at the direction of an enterprise, for example. Alternatively, in some embodiments, the training parameters are automatically selected based upon data provided in one or more training data sets **1306**. The training parameters can include, for example, a learning rate, a model size, a number of training passes, data shuffling, regularization, and/or other training parameters known to those skilled in the art. The training data in the training data sets **1306**.

(98) The learning rate is a training parameter defined by a constant value. The learning rate affects the speed at which the machine learning algorithm **1304** converges to the optimal weights. The machine learning algorithm **1304** can update the weights for every data example included in the training data set **1306**. The size of an update is controlled by the learning rate. A learning rate that is too high might prevent the machine learning algorithm **1304** from converging to the optimal weights. A learning rate that is too low might result in the machine learning algorithm **1304** requiring multiple training passes to converge to the optimal weights.

(99) The model size is regulated by the number of input features (“features”) **1308** in the training data set **1306**. A greater the number of features **1308** yields a greater number of possible patterns that can be determined from the training data set **1306**. The model size should be selected to balance the resources (e.g., compute, memory, storage, etc.) needed for training and the predictive power of the resultant machine learning model **1302**.

(100) The number of training passes indicates the number of training passes that the machine learning algorithm **1304** makes over the training data set **1306** during the training process. The number of training passes can be adjusted based, for example, on the size of the training data set **1306**, with larger training data sets being exposed to fewer training passes in consideration of time and/or resource utilization. The effectiveness of the resultant machine learning model **1302** can be increased by multiple training passes.

(101) Data shuffling is a training parameter designed to prevent the machine learning algorithm **1304** from reaching false optimal weights due to the order in which data contained in the training data set **1306** is processed. For example, data provided in rows and columns might be analyzed first row, second row, third row, etc., and thus an optimal weight might be obtained well before a full range of data has been considered. By data shuffling, the data contained in the training data set **1306** can be analyzed more thoroughly and mitigate bias in the resultant machine learning model **1302**.

(102) Regularization is a training parameter that helps to prevent the machine learning model **1302** from memorizing training data from the training data set **1306**. In other words, the machine learning model **1302** fits the training data set **1306**, but the predictive performance of the machine learning model **1302** is not acceptable. Regularization helps the machine learning system **1300** avoid this overfitting/memorization problem by adjusting extreme weight values of the features **1308**. For example, a feature that has a small weight value relative to the weight values of the other features in the training data set **1306** can be adjusted to zero.

(103) The machine learning system **1300** can determine model accuracy after training by using one

or more evaluation data sets **1310** containing the same features **1308'** as the features **1308** in the training data set **1306**. This also prevents the machine learning model **1302** from simply memorizing the data contained in the training data set **1306**. The number of evaluation passes made by the machine learning system **1300** can be regulated by a target model accuracy that, when reached, ends the evaluation process and the machine learning model **1302** is considered ready for deployment.

(104) After deployment, the machine learning model **1302** can perform a prediction operation (“prediction”) **1314** with an input data set **1312** having the same features **1308'** as the features **1308** in the training data set **1306** and the features **1308'** of the evaluation data set **1310**. The results of the prediction **1314** are included in an output data set **1316** consisting of predicted data. The machine learning model **1302** can perform other operations, such as regression, classification, and others. As such, the example illustrated in FIG. **13** should not be construed as being limiting in any way.

(105) Turning now to FIG. **14**, details of a network **1400** are illustrated, according to an illustrative embodiment. The network **1400** includes a cellular network **1402**, a packet data network **1404**, and a circuit switched network **1406**. In some embodiments, the network **818** is or includes the network **1400**. Moreover, the ICAM system **100** and/or the AAM system **600** can be configured to communicate over the network **1400**.

(106) The cellular network **1402** can include various components such as, but not limited to, base transceiver stations (“BTSs”), Node-Bs or e-Node-Bs, base station controllers (“BSCs”), radio network controllers (“RNCs”), mobile switching centers (“MSCs”), mobility management entities (“MMEs”), short message service centers (“SMSCs”), multimedia messaging service centers (“MMSCs”), home location registers (“HLRs”), home subscriber servers (“HS Ss”), visitor location registers (“VLRs”), charging platforms, billing platforms, voicemail platforms, GPRS core network components, location service nodes, and the like. The cellular network **1402** also includes radios and nodes for receiving and transmitting voice, data, and combinations thereof to and from radio transceivers, networks, the packet data network **1404**, and the circuit switched network **1406**.

(107) A mobile communications device **1408**, such as, for example, a cellular telephone, a user equipment, a mobile terminal, a PDA, a laptop computer, a handheld computer, and combinations thereof, can be operatively connected to the cellular network **1402**. The cellular network **1402** can be configured as a GSM network and can provide data communications via GPRS and/or EDGE. Additionally, or alternatively, the cellular network **1402** can be configured as a 3G Universal Mobile Telecommunications System (“UMTS”) network and can provide data communications via the HSPA protocol family, for example, HSDPA, EUL, and HSPA+. The cellular network **1402** also is compatible with 4G mobile communications standards such as LTE, 5G mobile communications standards, or the like, as well as evolved and future mobile standards.

(108) The packet data network **1404** includes various systems, devices, servers, computers, databases, and other devices in communication with one another, as is generally known. In some embodiments, the packet data network **1404** is or includes one or more WI-FI networks, each of which can include one or more WI-FI access points, routers, switches, and other WI-FI network components. The packet data network **1404** devices are accessible via one or more network links. The servers often store various files that are provided to a requesting device such as, for example, a computer, a terminal, a smartphone, or the like. Typically, the requesting device includes software for executing a web page in a format readable by the browser or other software. Other files and/or data may be accessible via “links” in the retrieved files, as is generally known. In some embodiments, the packet data network **1404** includes or is in communication with the Internet. The circuit switched network **1406** includes various hardware and software for providing circuit switched communications. The circuit switched network **1406** may include, or may be, what is often referred to as a plain old telephone system (“POTS”). The functionality of a circuit switched network **1406** or other circuit-switched network are generally known and will not be described

herein in detail.

(109) The illustrated cellular network **1402** is shown in communication with the packet data network **1404** and a circuit switched network **1406**, though it should be appreciated that this is not necessarily the case. One or more Internet-capable devices **1408** such as a laptop, a portable device, or another suitable device, can communicate with one or more cellular networks **1402**, and devices connected thereto, through the packet data network **1404**. It also should be appreciated that the Internet-capable device **1410** can communicate with the packet data network **1404** through the circuit switched network **1406**, the cellular network **1402**, and/or via other networks (not illustrated).

(110) As illustrated, a communications device **1412**, for example, a telephone, facsimile machine, modem, computer, or the like, can be in communication with the circuit switched network **1406**, and therethrough to the packet data network **1404** and/or the cellular network **1402**. It should be appreciated that the communications device **1412** can be an Internet-capable device, and can be substantially similar to the Internet-capable device **1410**.

(111) Based on the foregoing, it should be appreciated that concepts and technologies directed to image classification attack mitigation have been disclosed herein. Although the subject matter presented herein has been described in language specific to computer structural features, methodological and transformative acts, specific computing machinery, and computer-readable media, it is to be understood that the concepts and technologies disclosed herein are not necessarily limited to the specific features, acts, or media described herein. Rather, the specific features, acts and mediums are disclosed as example forms of implementing the concepts and technologies disclosed herein.

(112) The subject matter described above is provided by way of illustration only and should not be construed as limiting. Various modifications and changes may be made to the subject matter described herein without following the example embodiments and applications illustrated and described, and without departing from the true spirit and scope of the embodiments of the concepts and technologies disclosed herein.

## Claims

1. A method, comprising: processing, by a computing system, an original audio signal in a first fashion to generate a first digital audio signal; processing, by the computing system, the original audio signal in a second fashion, different than the first fashion, to generate a second digital audio signal different than the first digital audio signal; converting, by the computing system, the first digital audio signal and the second digital audio signal into first text and second text, respectively; determining, by the computing system, a mismatch between the first text and the second text; and based at least in part on the mismatch, sending an alarm indication from the computing system to a remote device.
2. The method of claim 1, wherein: processing the original audio signal in the first fashion includes sampling the original audio signal at a first bit depth to generate the first digital audio signal; and processing the original audio signal in the second fashion includes sampling the original audio signal at a second bit depth, different than the first bit depth, to generate the second digital audio signal.
3. The method of claim 1, wherein: processing the original audio signal in the first fashion includes sampling the original audio signal at a first sampling rate to generate the first digital audio signal; and processing the original audio signal in the second fashion includes sampling the original audio signal at a second sampling rate, different than the first sampling rate, to generate the second digital audio signal.
4. The method of claim 1, wherein: processing the original audio signal in the first fashion includes sampling the original audio signal to generate the first digital audio signal; and processing the

original audio signal in the second fashion includes enhancing the original audio signal to generate the second digital audio signal.

5. The method of claim 1, wherein: processing the original audio signal in the first fashion includes sampling the original audio signal to generate the first digital audio signal; and processing the original audio signal in the second fashion includes compressing the original audio signal to generate the second digital audio signal.

6. The method of claim 1, wherein: processing the original audio signal in the first fashion includes enhancing the original audio signal to generate the first digital audio signal; and processing the original audio signal in the second fashion includes compressing the original audio signal to generate the second digital audio signal.

7. The method of claim 1, further comprising: receiving the original audio signal by a voice-enabled system configured to control an operation of a home security component.

8. The method of claim 7, wherein: the home security component comprises a smart lock; and the voice-enabled system is configured to control operation of the smart lock in response to voice commands.

9. A system, comprising: one or more processors; and one or more non-transitory computer-readable mediums encoded with instructions which, when executed by the one or more processors, cause the system to: process an original audio signal in a first fashion to generate a first digital audio signal; process the original audio signal in a second fashion, different than the first fashion, to generate a second digital audio signal different than the first digital audio signal; convert the first digital audio signal and the second digital audio signal into first text and second text, respectively; determine a mismatch between the first text and the second text; and based at least in part on the mismatch, send an alarm indication to a remote device.

10. The system of claim 9, wherein the one or more non-transitory computer-readable mediums are further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by sampling the original audio signal at a first bit depth to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by sampling the original audio signal at a second bit depth, different than the first bit depth, to generate the second digital audio signal.

11. The system of claim 9, wherein the one or more non-transitory computer-readable mediums are further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by sampling the original audio signal at a first sampling rate to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by sampling the original audio signal at a second sampling rate, different than the first sampling rate, to generate the second digital audio signal.

12. The system of claim 9, wherein the one or more non-transitory computer-readable mediums are further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by sampling the original audio signal to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by enhancing the original audio signal to generate the second digital audio signal.

13. The system of claim 9, wherein the one or more non-transitory computer-readable mediums are further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by sampling the original audio signal to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by compressing the original audio signal to generate the second digital audio signal.

14. The system of claim 9, wherein the one or more non-transitory computer-readable mediums are

further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by enhancing the original audio signal to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by compressing the original audio signal to generate the second digital audio signal.

15. The system of claim 9, wherein the one or more non-transitory computer-readable mediums are further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: control an operation of a home security component in response to receipt of the original audio signal.

16. The system of claim 15, wherein the home security component comprises a smart lock.

17. One or more non-transitory computer-readable mediums encoded with instructions which, when executed by one or more processors of a system, cause the system to: process an original audio signal in a first fashion to generate a first digital audio signal; process the original audio signal in a second fashion, different than the first fashion, to generate a second digital audio signal different than the first digital audio signal; convert the first digital audio signal and the second digital audio signal into first text and second text, respectively; determine a mismatch between the first text and the second text; and based at least in part on the mismatch, send an alarm indication to a remote device.

18. The one or more non-transitory computer-readable mediums of claim 17, further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by sampling the original audio signal at a first bit depth to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by sampling the original audio signal at a second bit depth, different than the first bit depth, to generate the second digital audio signal.

19. The one or more non-transitory computer-readable mediums of claim 17, further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by sampling the original audio signal at a first sampling rate to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by sampling the original audio signal at a second sampling rate, different than the first sampling rate, to generate the second digital audio signal.

20. The one or more non-transitory computer-readable mediums of claim 17, further encoded with additional instructions which, when executed by the one or more processors, further cause the system to: process the original audio signal in the first fashion at least in part by enhancing the original audio signal to generate the first digital audio signal; and process the original audio signal in the second fashion at least in part by compressing the original audio signal to generate the second digital audio signal.

---