

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393677
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Sutherland; Andrew et al.

Systems and methods for metadata classification

Abstract

Systems and methods for metadata processing. In some embodiments, one or more metadata inputs may be processed to determine whether to allow an instruction. For instance, one or more classification bits may be identified from a metadata input of the one or more metadata inputs, and the metadata input may be processed based on the one or more classification bits.

Inventors:	Sutherland; Andrew (Portland, OR), Milburn; Steven (Cranston, RI), Sullivan; Gregory T. (Auburndale, MA), Boling; Eli (Manchester, MA)
Applicant:	Dover Microsystems, Inc. (Wayland, MA)
Family ID:	1000008764127
Assignee:	Dover Microsystems, Inc. (Wayland, MA)
Appl. No.:	17/423701
Filed (or PCT Filed):	January 15, 2020
PCT No.:	PCT/US2020/013678
PCT Pub. No.:	WO2020/150351
PCT Pub. Date:	July 23, 2020

Prior Publication Data

Document Identifier	Publication Date
US 20220092173 A1	Mar. 24, 2022

Related U.S. Application Data

us-provisional-application US 62794499 20190118

Publication Classification

Int. Cl.: G06F21/54 (20130101); G06F21/12 (20130101); G06F21/51 (20130101); G06F21/62 (20130101); G06F21/78 (20130101)

U.S. Cl.:

CPC G06F21/54 (20130101); G06F21/125 (20130101); G06F21/51 (20130101); G06F21/6209 (20130101); G06F21/78 (20130101);

Field of Classification Search

CPC: G06F (21/54); G06F (21/125); G06F (21/51); G06F (21/6209); G06F (21/78)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5377336	12/1993	Eickemeyer et al.	N/A	N/A
5488730	12/1995	Brown, III et al.	N/A	N/A
5521910	12/1995	Matthews	N/A	N/A
5559952	12/1995	Fujimoto	N/A	N/A
5655100	12/1996	Ebrahim et al.	N/A	N/A
5664197	12/1996	Kardach et al.	N/A	N/A
5664223	12/1996	Bender et al.	N/A	N/A
5684977	12/1996	Van Loo et al.	N/A	N/A
5764946	12/1997	Tran et al.	N/A	N/A
5778423	12/1997	Sites et al.	N/A	N/A
5848433	12/1997	Tran et al.	N/A	N/A
5890008	12/1998	Panwar et al.	N/A	N/A
5941981	12/1998	Tran	N/A	N/A
6014734	12/1999	Tran et al.	N/A	N/A
6035374	12/1999	Panwar et al.	N/A	N/A
6058466	12/1999	Panwar et al.	N/A	N/A
6240502	12/2000	Panwar et al.	N/A	N/A
6269436	12/2000	Tran et al.	N/A	N/A
6298432	12/2000	Goto	N/A	N/A
6321297	12/2000	Shamanna et al.	N/A	N/A
6324599	12/2000	Zhou et al.	N/A	N/A
6351784	12/2001	Neal et al.	N/A	N/A
6438673	12/2001	Jourdan et al.	N/A	N/A
6449714	12/2001	Sinharoy	N/A	N/A
6549903	12/2002	Babaian et al.	N/A	N/A
6549959	12/2002	Yates et al.	N/A	N/A
6622182	12/2002	Miller et al.	N/A	N/A
6625146	12/2002	Merchant et al.	N/A	N/A
6636523	12/2002	Lau et al.	N/A	N/A
6748589	12/2003	Johnson et al.	N/A	N/A
6922740	12/2004	Kondratiev et al.	N/A	N/A

6976147	12/2004	Isaac et al.	N/A	N/A
7047394	12/2005	Van Dyke et al.	N/A	N/A
7058918	12/2005	Abramovici et al.	N/A	N/A
7095783	12/2005	Sotheran et al.	N/A	N/A
7137086	12/2005	Abramovici	N/A	N/A
7146548	12/2005	Abramovici	N/A	N/A
7213247	12/2006	Wilner et al.	N/A	N/A
7263572	12/2006	Hiji	N/A	N/A
7296201	12/2006	Abramovici	N/A	N/A
7301541	12/2006	Hansen et al.	N/A	N/A
7305635	12/2006	Abramovici et al.	N/A	N/A
7313820	12/2006	Kilroy	N/A	N/A
7340469	12/2007	Alghathbar et al.	N/A	N/A
7348796	12/2007	Crouch et al.	N/A	N/A
7426705	12/2007	Kolaric	N/A	N/A
7430650	12/2007	Ross	N/A	N/A
7434002	12/2007	Zedlewski et al.	N/A	N/A
7467414	12/2007	Schlesinger	N/A	N/A
7487264	12/2008	Pandya	N/A	N/A
7493247	12/2008	Memmi	N/A	N/A
7493434	12/2008	Abramovici	N/A	N/A
7574536	12/2008	Johnsen et al.	N/A	N/A
7580914	12/2008	Wang et al.	N/A	N/A
7581064	12/2008	Zedlewski et al.	N/A	N/A
7631107	12/2008	Pandya	N/A	N/A
7688838	12/2009	Aloni et al.	N/A	N/A
7698402	12/2009	Santos et al.	N/A	N/A
7813342	12/2009	Gadelrab	N/A	N/A
7886148	12/2010	Kiriansky et al.	N/A	N/A
8074052	12/2010	Iadonato et al.	N/A	N/A
8121828	12/2011	Yates, Jr. et al.	N/A	N/A
8127121	12/2011	Yates, Jr. et al.	N/A	N/A
8131663	12/2011	Taylor	N/A	N/A
8131762	12/2011	Smetters et al.	N/A	N/A
8166404	12/2011	Grinstein	N/A	N/A
8181219	12/2011	Golan et al.	N/A	N/A
8271447	12/2011	Natanzon et al.	N/A	N/A
8326774	12/2011	Candelore	N/A	N/A
8335754	12/2011	Dawson et al.	N/A	N/A
8346760	12/2012	Wang et al.	N/A	N/A
8364910	12/2012	Wilkerson et al.	N/A	N/A
8380933	12/2012	Uchiyama	N/A	N/A
8423720	12/2012	Döring	N/A	N/A
8516193	12/2012	Clinton et al.	N/A	N/A
8543606	12/2012	Pulfer et al.	N/A	N/A
8572410	12/2012	Tkacik et al.	N/A	N/A
8677457	12/2013	Fullerton	N/A	N/A
8701200	12/2013	Naldurg et al.	N/A	N/A
8738860	12/2013	Griffin et al.	N/A	N/A
8756185	12/2013	Dawson et al.	N/A	N/A

8788792	12/2013	Yates, Jr. et al.	N/A	N/A
8826391	12/2013	Tkacik et al.	N/A	N/A
8843734	12/2013	Lim	N/A	N/A
8874850	12/2013	Goodson et al.	N/A	N/A
8875170	12/2013	Daniel	N/A	N/A
8887230	12/2013	Barton et al.	N/A	N/A
8904477	12/2013	Barton et al.	N/A	N/A
8938783	12/2014	Becker et al.	N/A	N/A
9026733	12/2014	Clinton et al.	N/A	N/A
9047178	12/2014	Talagala et al.	N/A	N/A
9087200	12/2014	McKeen et al.	N/A	N/A
9165078	12/2014	Bester et al.	N/A	N/A
9208082	12/2014	Cheriton et al.	N/A	N/A
9219752	12/2014	Balinsky et al.	N/A	N/A
9251052	12/2015	Talagala et al.	N/A	N/A
9311093	12/2015	Gschwind et al.	N/A	N/A
9317708	12/2015	Lee et al.	N/A	N/A
9323684	12/2015	Koker et al.	N/A	N/A
9467474	12/2015	Barton et al.	N/A	N/A
9507589	12/2015	Rao et al.	N/A	N/A
9507598	12/2015	Bonnano et al.	N/A	N/A
9513884	12/2015	Bates et al.	N/A	N/A
9525606	12/2015	Staggs et al.	N/A	N/A
9571509	12/2016	Satish et al.	N/A	N/A
9576147	12/2016	McClintock et al.	N/A	N/A
9665603	12/2016	Bester et al.	N/A	N/A
9680736	12/2016	Kamboh	N/A	N/A
9680738	12/2016	Jackson et al.	N/A	N/A
9703956	12/2016	Watson et al.	N/A	N/A
9736185	12/2016	Belamaric et al.	N/A	N/A
9785440	12/2016	DeHon	N/A	N/A
9792472	12/2016	Robshaw et al.	N/A	N/A
9906557	12/2017	Hsiung et al.	N/A	N/A
9953095	12/2017	Scott et al.	N/A	N/A
10073977	12/2017	Pappachan et al.	N/A	N/A
10078763	12/2017	Chiricescu et al.	N/A	N/A
10114958	12/2017	Sell	N/A	N/A
10133866	12/2017	Kumar et al.	N/A	N/A
10152330	12/2017	Chiricescu et al.	N/A	N/A
10235176	12/2018	DeHon et al.	N/A	N/A
10261794	12/2018	DeHon	N/A	N/A
10424043	12/2018	Koston et al.	N/A	N/A
10503904	12/2018	Singh et al.	N/A	N/A
10521230	12/2018	DeHon	N/A	N/A
10545760	12/2019	DeHon	N/A	N/A
10642616	12/2019	DeHon et al.	N/A	N/A
10642753	12/2019	Steinberg	N/A	N/A
10719630	12/2019	Chiricescu et al.	N/A	N/A
10725778	12/2019	DeHon et al.	N/A	N/A
10754650	12/2019	DeHon et al.	N/A	N/A

10936713	12/2020	DeHon et al.	N/A	N/A
11150910	12/2020	Milburn et al.	N/A	N/A
11182162	12/2020	DeHon et al.	N/A	N/A
11340902	12/2021	DeHon	N/A	N/A
11417109	12/2021	Theimer et al.	N/A	N/A
11507373	12/2021	Dehon et al.	N/A	N/A
11556664	12/2022	Levy et al.	N/A	N/A
11635960	12/2022	DeHon	N/A	N/A
11841956	12/2022	Sullivan	N/A	G06F 21/602
12079197	12/2023	Boling	N/A	G06F 9/30101
2002/0083298	12/2001	Cook et al.	N/A	N/A
2002/0087795	12/2001	Hum et al.	N/A	N/A
2002/0124156	12/2001	Yoaz et al.	N/A	N/A
2003/0014466	12/2002	Berger et al.	N/A	N/A
2003/0023783	12/2002	Arimilli et al.	N/A	N/A
2003/0058889	12/2002	Lansing et al.	N/A	N/A
2003/0120892	12/2002	Hum et al.	N/A	N/A
2003/0145235	12/2002	Choo	N/A	N/A
2003/0149895	12/2002	Choo et al.	N/A	N/A
2003/0172109	12/2002	Dalton et al.	N/A	N/A
2003/0196108	12/2002	Kung	N/A	N/A
2004/0015845	12/2003	Hickman	N/A	N/A
2004/0117599	12/2003	Mittal et al.	N/A	N/A
2004/0133777	12/2003	Kiriansky et al.	N/A	N/A
2004/0236876	12/2003	Kondratiev et al.	N/A	N/A
2004/0255267	12/2003	Meijer	N/A	N/A
2005/0055565	12/2004	Fournet et al.	N/A	N/A
2005/0108518	12/2004	Pandya	N/A	N/A
2005/0149521	12/2004	Wang et al.	N/A	N/A
2005/0149719	12/2004	Kilroy	N/A	N/A
2005/0154838	12/2004	DeWitt et al.	N/A	N/A
2005/0182667	12/2004	Metzger et al.	N/A	N/A
2006/0004548	12/2005	Santos et al.	N/A	N/A
2006/0059567	12/2005	Bird et al.	N/A	N/A
2006/0080489	12/2005	Hiji	N/A	N/A
2006/0080553	12/2005	Hall	N/A	N/A
2006/0090084	12/2005	Buer	N/A	N/A
2006/0112261	12/2005	Yourst et al.	N/A	N/A
2006/0143689	12/2005	Yu et al.	N/A	N/A
2006/0242332	12/2005	Johnsen et al.	N/A	N/A
2006/0277392	12/2005	Bittner	N/A	N/A
2007/0006294	12/2006	Hunter	N/A	N/A
2007/0050586	12/2006	Shin et al.	N/A	N/A
2007/0226365	12/2006	Hildreth et al.	N/A	N/A
2007/0239861	12/2006	Reeves et al.	N/A	N/A
2007/0261033	12/2006	Chen et al.	N/A	N/A
2007/0279264	12/2006	Nakagawa	N/A	N/A
2008/0010233	12/2007	Sack et al.	N/A	N/A
2008/0016547	12/2007	Anderson et al.	N/A	N/A

2008/0028196	12/2007	Nailas	N/A	N/A
2008/0052488	12/2007	Fritz et al.	N/A	N/A
2008/0066160	12/2007	Becker et al.	N/A	N/A
2008/0083298	12/2007	Lin	N/A	N/A
2008/0126841	12/2007	Benhanokh et al.	N/A	N/A
2008/0140737	12/2007	Garst et al.	N/A	N/A
2008/0168529	12/2007	Anderson et al.	N/A	N/A
2008/0201333	12/2007	Rowley	N/A	N/A
2008/0216073	12/2007	Yates et al.	N/A	N/A
2008/0216102	12/2007	Quinn	N/A	N/A
2008/0222397	12/2007	Wilkerson et al.	N/A	N/A
2008/0240111	12/2007	Gadelrab	N/A	N/A
2008/0244232	12/2007	Sherman et al.	N/A	N/A
2008/0248599	12/2007	Jaiswal et al.	N/A	N/A
2008/0250216	12/2007	Kershaw et al.	N/A	N/A
2008/0282040	12/2007	Doring	N/A	N/A
2008/0282093	12/2007	Hatakeyama	N/A	N/A
2008/0288941	12/2007	Adams et al.	N/A	N/A
2008/0301256	12/2007	McWilliams et al.	N/A	N/A
2008/0301471	12/2007	Demarest et al.	N/A	N/A
2008/0320235	12/2007	Beckmann et al.	N/A	N/A
2009/0006519	12/2008	Nandan et al.	N/A	N/A
2009/0063951	12/2008	Rjaibi et al.	N/A	N/A
2009/0097815	12/2008	Lahr et al.	N/A	N/A
2009/0113110	12/2008	Chen et al.	N/A	N/A
2009/0113132	12/2008	Cain, III et al.	N/A	N/A
2009/0113135	12/2008	Cain et al.	N/A	N/A
2009/0144388	12/2008	Gross et al.	N/A	N/A
2009/0164705	12/2008	Gorobets	N/A	N/A
2009/0164766	12/2008	Suggs et al.	N/A	N/A
2009/0165078	12/2008	Samudrala et al.	N/A	N/A
2009/0178102	12/2008	Alghathbar et al.	N/A	N/A
2009/0204785	12/2008	Yates, Jr. et al.	N/A	N/A
2009/0241097	12/2008	Wang et al.	N/A	N/A
2009/0254543	12/2008	Ber et al.	N/A	N/A
2009/0254572	12/2008	Redlich et al.	N/A	N/A
2010/0011209	12/2009	Kiriansky et al.	N/A	N/A
2010/0011446	12/2009	Klucher et al.	N/A	N/A
2010/0022869	12/2009	Kimura	N/A	N/A
2010/0049974	12/2009	Winjum et al.	N/A	N/A
2010/0138613	12/2009	Parker	N/A	N/A
2010/0154026	12/2009	Chatterjee et al.	N/A	N/A
2010/0169382	12/2009	Sheaffer et al.	N/A	N/A
2010/0191922	12/2009	Dickey et al.	N/A	N/A
2010/0228693	12/2009	Dawson et al.	N/A	N/A
2010/0235580	12/2009	Bouvier	N/A	N/A
2010/0250729	12/2009	Morris	N/A	N/A
2010/0332456	12/2009	Prahlad et al.	N/A	N/A
2010/0332716	12/2009	Sheaffer et al.	N/A	N/A
2011/0016295	12/2010	Catherwood et al.	N/A	N/A

2011/0078389	12/2010	Patel et al.	N/A	N/A
2011/0099336	12/2010	Yasufuku et al.	N/A	N/A
2011/0126265	12/2010	Fullerton	N/A	N/A
2011/0161623	12/2010	Eichenberger et al.	N/A	N/A
2011/0219424	12/2010	Panasyuk et al.	N/A	N/A
2011/0238805	12/2010	Signori	N/A	N/A
2012/0036507	12/2011	Jonnala et al.	N/A	N/A
2012/0079458	12/2011	Williams et al.	N/A	N/A
2012/0117610	12/2011	Pandya	N/A	N/A
2012/0144167	12/2011	Yates, Jr. et al.	N/A	N/A
2012/0151184	12/2011	Wilkerson et al.	N/A	N/A
2012/0180031	12/2011	Eichenberger et al.	N/A	N/A
2012/0203970	12/2011	Biran et al.	N/A	N/A
2012/0210066	12/2011	Joshi et al.	N/A	N/A
2012/0210068	12/2011	Joshi et al.	N/A	N/A
2012/0233212	12/2011	Newton et al.	N/A	N/A
2012/0236756	12/2011	Bennett et al.	N/A	N/A
2012/0255018	12/2011	Sallam	N/A	N/A
2012/0297057	12/2011	Ghosh et al.	N/A	N/A
2013/0006993	12/2012	Kobayashi	N/A	N/A
2013/0010779	12/2012	Fischer et al.	N/A	N/A
2013/0016075	12/2012	Kim et al.	N/A	N/A
2013/0047142	12/2012	Bates et al.	N/A	N/A
2013/0067593	12/2012	Candelore	N/A	N/A
2013/0081134	12/2012	Glew et al.	N/A	N/A
2013/0097369	12/2012	Talagala et al.	N/A	N/A
2013/0097421	12/2012	Lim	N/A	N/A
2013/0097667	12/2012	Pulfer et al.	N/A	N/A
2013/0138892	12/2012	Loh et al.	N/A	N/A
2013/0159726	12/2012	McKeen et al.	N/A	N/A
2013/0160075	12/2012	Schlesinger et al.	N/A	N/A
2013/0160775	12/2012	Curnow	N/A	N/A
2013/0185475	12/2012	Talagala et al.	N/A	N/A
2013/0185488	12/2012	Talagala et al.	N/A	N/A
2013/0212321	12/2012	Talagala et al.	N/A	N/A
2013/0227218	12/2012	Chang et al.	N/A	N/A
2013/0254838	12/2012	Ahuja et al.	N/A	N/A
2013/0275656	12/2012	Talagala et al.	N/A	N/A
2013/0283017	12/2012	Wilkerson et al.	N/A	N/A
2013/0290607	12/2012	Chang et al.	N/A	N/A
2013/0312099	12/2012	Edwards et al.	N/A	N/A
2013/0326117	12/2012	Aune	N/A	N/A
2014/0006804	12/2013	Tkacik et al.	N/A	N/A
2014/0019385	12/2013	Dawson et al.	N/A	N/A
2014/0047181	12/2013	Peterson et al.	N/A	N/A
2014/0047183	12/2013	Chawla et al.	N/A	N/A
2014/0101396	12/2013	Bonanno et al.	N/A	N/A
2014/0140342	12/2013	Narad	N/A	N/A

2014/0173211	12/2013	Loh et al.	N/A	N/A
2014/0223445	12/2013	Beckmann et al.	N/A	N/A
2014/0280248	12/2013	Bester et al.	N/A	N/A
2014/0281192	12/2013	Gilda et al.	N/A	N/A
2014/0282832	12/2013	Chanoch et al.	N/A	N/A
2014/0283040	12/2013	Wilkerson et al.	N/A	N/A
2014/0283107	12/2013	Walton et al.	N/A	N/A
2015/0012689	12/2014	Atkisson et al.	N/A	N/A
2015/0046658	12/2014	Wilson	N/A	N/A
2015/0058997	12/2014	Lee et al.	N/A	N/A
2015/0089186	12/2014	Kim et al.	N/A	N/A
2015/0092778	12/2014	Jackson et al.	N/A	N/A
2015/0120699	12/2014	Faerber et al.	N/A	N/A
2015/0149673	12/2014	Balkan et al.	N/A	N/A
2015/0205535	12/2014	Joshi et al.	N/A	N/A
2015/0220453	12/2014	Heisswolf et al.	N/A	N/A
2015/0249668	12/2014	Reddy et al.	N/A	N/A
2015/0278311	12/2014	Isherwood et al.	N/A	N/A
2015/0339062	12/2014	Toyoda et al.	N/A	N/A
2015/0339329	12/2014	Bester et al.	N/A	N/A
2015/0378780	12/2014	Busaba et al.	N/A	N/A
2015/0381660	12/2014	Hsiung et al.	N/A	N/A
2016/0048551	12/2015	Baldwin et al.	N/A	N/A
2016/0062803	12/2015	Beckmann et al.	N/A	N/A
2016/0077816	12/2015	Eilam et al.	N/A	N/A
2016/0092702	12/2015	Durham et al.	N/A	N/A
2016/0132536	12/2015	Lee	N/A	N/A
2016/0140363	12/2015	Chiricescu et al.	N/A	N/A
2016/0154833	12/2015	Isherwood, Jr. et al.	N/A	N/A
2016/0170769	12/2015	LeMay	N/A	N/A
2016/0182408	12/2015	Jani et al.	N/A	N/A
2016/0188891	12/2015	Schlesinger et al.	N/A	N/A
2016/0196432	12/2015	Main et al.	N/A	N/A
2016/0239223	12/2015	Joshi et al.	N/A	N/A
2016/0274810	12/2015	Godard et al.	N/A	N/A
2016/0335187	12/2015	Greenspan et al.	N/A	N/A
2016/0350019	12/2015	Koufaty et al.	N/A	N/A
2016/0350230	12/2015	Murphy	N/A	N/A
2016/0366102	12/2015	Smith	N/A	N/A
2016/0371496	12/2015	Sell	N/A	N/A
2017/0024568	12/2016	Pappachan et al.	N/A	N/A
2017/0031708	12/2016	Chen et al.	N/A	N/A
2017/0048249	12/2016	Berrangé	N/A	N/A
2017/0061160	12/2016	Simonov et al.	N/A	N/A
2017/0083338	12/2016	Burger et al.	N/A	N/A
2017/0091107	12/2016	Peterson et al.	N/A	N/A
2017/0126687	12/2016	Martinelli	N/A	N/A
2017/0126738	12/2016	Wilkerson et al.	N/A	N/A
2017/0177367	12/2016	DeHon	N/A	N/A

2017/0177368	12/2016	DeHon et al.	N/A	N/A
2017/0192986	12/2016	Isherwood et al.	N/A	N/A
2017/0220806	12/2016	Munoz et al.	N/A	N/A
2017/0235840	12/2016	Bester et al.	N/A	N/A
2017/0286119	12/2016	Al Sheikh et al.	N/A	N/A
2017/0286151	12/2016	Landers et al.	N/A	N/A
2017/0293563	12/2016	DeHon et al.	N/A	N/A
2017/0308480	12/2016	Wilson	N/A	N/A
2017/0329961	12/2016	Shanbhogue et al.	N/A	N/A
2018/0011708	12/2017	DeHon	N/A	N/A
2018/0046579	12/2017	Greenspan et al.	N/A	N/A
2018/0082055	12/2017	Fleming et al.	N/A	N/A
2018/0121650	12/2017	Brown	N/A	N/A
2018/0143890	12/2017	Ogawa et al.	N/A	N/A
2018/0189062	12/2017	Baghsorkhi et al.	N/A	N/A
2018/0276085	12/2017	Mitkar et al.	N/A	N/A
2018/0302443	12/2017	Weiss et al.	N/A	N/A
2018/0336031	12/2017	DeHon et al.	N/A	N/A
2018/0336032	12/2017	DeHon et al.	N/A	N/A
2018/0336033	12/2017	DeHon	N/A	N/A
2018/0341490	12/2017	DeHon	N/A	N/A
2019/0034665	12/2018	Chiricescu et al.	N/A	N/A
2019/0141059	12/2018	Shimizu et al.	N/A	N/A
2019/0155606	12/2018	Rotem et al.	N/A	N/A
2019/0171457	12/2018	DeHon et al.	N/A	N/A
2019/0205244	12/2018	Smith	N/A	N/A
2019/0213322	12/2018	DeHon et al.	N/A	N/A
2019/0236272	12/2018	Piatt	N/A	N/A
2019/0243655	12/2018	Milburn et al.	N/A	N/A
2019/0243768	12/2018	Doshi et al.	N/A	N/A
2019/0354675	12/2018	Gan et al.	N/A	N/A
2019/0384604	12/2018	DeHon et al.	N/A	N/A
2019/0392146	12/2018	Gezalov et al.	N/A	N/A
2019/0392147	12/2018	Gezalov et al.	N/A	N/A
2020/0089500	12/2019	DeHon	N/A	N/A
2020/0125502	12/2019	Durham et al.	N/A	N/A
2020/0201576	12/2019	Yudanov et al.	N/A	N/A
2020/0387374	12/2019	DeHon	N/A	G06F 9/3867
2020/0387384	12/2019	Huang	N/A	N/A
2021/0004231	12/2020	DeHon	N/A	N/A
2021/0026934	12/2020	Boling et al.	N/A	N/A
2021/0042100	12/2020	Boling et al.	N/A	N/A
2021/0055954	12/2020	Milburn et al.	N/A	N/A
2021/0073375	12/2020	Milburn et al.	N/A	N/A
2021/0075797	12/2020	Gan et al.	N/A	N/A
2021/0255890	12/2020	Milburn et al.	N/A	N/A
2021/0406028	12/2020	Boling et al.	N/A	N/A
2021/0406137	12/2020	Sutherland et al.	N/A	N/A
2022/0012329	12/2021	Boling et al.	N/A	N/A
2022/0043654	12/2021	DeHon et al.	N/A	N/A

2022/0050904	12/2021	Sullivan et al.	N/A	N/A
2022/0129343	12/2021	Milburn et al.	N/A	N/A
2022/0198014	12/2021	Boling et al.	N/A	N/A
2022/0300583	12/2021	Boling et al.	N/A	N/A
2022/0309134	12/2021	Boling et al.	N/A	N/A
2022/0374415	12/2021	Boling et al.	N/A	N/A
2022/0398312	12/2021	Sutherland et al.	N/A	N/A
2023/0054942	12/2022	Milburn et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101558388	12/2008	CN	N/A
102160033	12/2010	CN	N/A
103282889	12/2012	CN	N/A
104657500	12/2014	CN	N/A
104794067	12/2014	CN	N/A
2519608	12/2014	GB	N/A
2013-242633	12/2012	JP	N/A
201333722	12/2012	TW	N/A
201729068	12/2016	TW	N/A
201737675	12/2016	TW	N/A
WO 92/03779	12/1991	WO	N/A
WO 2010/028316	12/2009	WO	N/A
WO 2015/047295	12/2014	WO	N/A
WO 2015/183481	12/2014	WO	N/A
WO 2017/106101	12/2016	WO	N/A
WO 2017/106103	12/2016	WO	N/A
WO 2017/221373	12/2016	WO	N/A
WO 2019/152772	12/2018	WO	N/A
WO 2019/152792	12/2018	WO	N/A
WO 2019/152795	12/2018	WO	N/A
WO 2019/152805	12/2018	WO	N/A
WO 2019/152822	12/2018	WO	N/A
WO 2019/213061	12/2018	WO	N/A
WO 2020/097177	12/2019	WO	N/A
WO 2020/097179	12/2019	WO	N/A
WO 2020/102064	12/2019	WO	N/A
WO 2020/132012	12/2019	WO	N/A
WO 2020/150351	12/2019	WO	N/A
WO 2021/076871	12/2020	WO	N/A
WO 2021/092138	12/2020	WO	N/A

OTHER PUBLICATIONS

International Search Report and Written Opinion for International Application No. PCT/US2022/042492 mailed Jan. 11, 2023. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/016242 mailed Mar. 29, 2019. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2016/066194 mailed Apr. 7, 2017. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2016/066188 mailed Jul. 13, 2017. cited by applicant

Invitation to Pay Additional Fees for International Application No. PCT/US2019/016295 mailed May 3, 2019. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/016295 mailed Jun. 26, 2019. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2019/016295 mailed Aug. 13, 2020. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/016276 mailed Apr. 26, 2019. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/029880 mailed Aug. 8, 2019. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/016317 mailed Apr. 16, 2019. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2020/059057 mailed Feb. 8, 2021. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/060030 mailed May 20, 2021. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/060698 mailed May 20, 2021. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/067084 mailed Jul. 1, 2021. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/013678 mailed Jul. 29, 2021. cited by applicant

Invitation to Pay Additional Fees for International Application No. PCT/US2019/013678 mailed Apr. 15, 2020. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2021/020602 mailed May 11, 2021. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2020/055952 mailed Jan. 28, 2021. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2019/060028 mailed May 20, 2021. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/060028 mailed Mar. 23, 2020. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/016272 mailed Apr. 8, 2019. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2019/016272 mailed Aug. 13, 2020. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2019/016276 mailed Aug. 13, 2020. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2019/029880 mailed Nov. 12, 2020. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2019/016317 mailed Aug. 13, 2020. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2020/055952 mailed Aug. 13, 2020. cited by applicant

[No Author Listed], Arm Limited: AMBA® AXI™ and ACE™ Protocol Specification. Oct. 28, 2011:1-306. [https://capocaccia.ethz.ch/capo/raw-attachment/wiki/2014/microblaze14/AX14_specification.pdf]. cited by applicant

Alves-Foss et al., Evaluating the Use of Security Tags in Security Policy Enforcement

Mechanisms. 2015 48th Hawaii International Conference on System Sciences Jan. 5, 2015:5201-10. cited by applicant

Berlekamp, Algebraic coding theory (revised edition). World Scientific; Mar. 26, 2015. 34 pages. cited by applicant

Calder et al., Process Algebra for Event-Driven Runtime Verification: A Case Study of Wireless Network Management. IFM LNCS 2012;21-23. cited by applicant

Dalton et al., Raksha: a flexible information flow architecture for software security. ACM Proceedings of the 34.SUP.th .Annual International Symposium In Computer Architecture. Jun. 9, 2007;35(2):482-93. cited by applicant

Dalton et al., Real-World Buffer Overflow Protection for Userspace and Kernelspace. USENIX Security Symposium Jul. 28, 2008;395-410. cited by applicant

De Amorim et al., A verified information-flow architecture. Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages Jan. 8, 2014:165-78. cited by applicant

De Amorim et al., Micro-policies: Formally verified, tag-based security monitors. 2015 IEEE Symposium on Security and Privacy May 17, 2015:813-30. cited by applicant

Dehon et al., Dover A Metadata-Extended RISC-V. Jan. 6, 2016:34.
[https://web.archive.org/web/20160331131339/http://riscv.org/wp-content/uploads/2016/01/Wed1430-dover_riscv_Jan2016_v3.pdf]. cited by applicant

Dehon et al., Dover: A metadata-extended Risc-V. Risc-V Workshop Oracle Conference Center, Redwood Shores, CA. Jan. 6, 2016. 33 pages. cited by applicant

Dhawan et al., Architectural support for software-defined metadata processing. Proceedings of the Twentieth International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS). Mar. 14, 2015:487-502. cited by applicant

Dhawan et al., Area-efficient near-associative memories on FPGAs. ACM Transactions on Reconfigurable Technology and Systems (TRETS). Jan. 23, 2015;7(4):1-22. cited by applicant

Dhawan et al., PUMP: a programmable unit for metadata processing. Proceedings of the Third Workshop on Hardware and Architectural Support for Security and Privacy (HASP). Jun. 15, 2014:1-8. cited by applicant

Engelke et al., Efficient LLVM-Based Dynamic Binary Translation. ACM VEE. Apr. 16, 2021;165-71. cited by applicant

Evans et al., Melding security metadata between software and hardware. Proceedings of the Posters and Demo Track Dec. 3, 2012:1-2. cited by applicant

Geater, Tee requirements for ISA. Thales eSecurity. 2018. 21 pages. cited by applicant

Hrițcu, Micro-policies: Formally verified, tag-based security monitors. Proceedings of the 10th ACM Workshop on Programming Languages and Analysis for Security Jul. 4, 2015. cited by applicant

Hunt et al., The Seven Properties of Highly Secure Devices. 2021. 10 pages. cited by applicant

Juglaret et al., Towards a fully abstract compiler using Micro-Policies: Secure compilation for mutually distrustful components. arXiv preprint arXiv:1510.00697. Oct. 2, 2015. 31 pages. cited by applicant

Kane, Runtime Monitoring for Safety-Critical Embedded Systems. Carnegie Mellon University. Feb. 2015;207 pages. cited by applicant

Kannan et al., Decoupling dynamic information flow tracking with a dedicated coprocessor. 2009 IEEE/IFIP International Conference on Dependable Systems & Networks. Jun. 29, 2009:105-14. cited by applicant

Mambretti et al., Trellis: Privilege separation for multi-user applications made easy. International Symposium on Research in Attacks, Intrusions, and Defenses. Springer, Cham. Sep. 19, 2016:437-56. cited by applicant

Nagarakatte et al., SoftBound: Highly Compatible and Complete Spatial Memory Safety for C.

University of Pennsylvania Department of Computer and Information Science Technical Report. Jan. 2009. 12 pages. cited by applicant

Okhravi et al., One Giant Leap for Computer Security. IEEE Computer and Reliability Societies. Jul./Aug. 2020;18(4):8-19. cited by applicant

Ozsoy et al., SIFT: A low-overhead dynamic information flow tracking architecture for smt processors. Proceedings of the 8th ACM International Conference on Computing Frontiers. May 3, 2011:1-11. cited by applicant

Roessler et al., Protecting the stack with metadata policies and tagged hardware. 2018 IEEE Symposium on Security and Privacy (SP) May 20, 2018;478-95. cited by applicant

Song et al., Security tagging for a zero-kernel operating system. 2013 46th Hawaii International Conference on System Sciences Jan. 7, 2013:5049-58. cited by applicant

Song et al., The 4th lowRISC Release: Tagged Memory and Minion Core. University of Cambridge. May 2017. 20 pages. https://riscv.org/wp-content/uploads/2017/05/Wed0930riscv201705_ppt.pdf [Last accessed Mar. 31, 2022]. cited by applicant

Suh et al., Secure Program Execution via Dynamic Information Flow Tracking. MIT CSAIL. Jul. 2003. 14 pages. cited by applicant

Sullivan et al., The dover inherently secure processor. 2017 IEEE International Symposium on Technologies for Homeland Security (HST) Apr. 25, 2017:1-5. cited by applicant

Tiwari et al., A small cache of large ranges: Hardware methods for efficiently searching, storing, and updating big dataflow tags. 2008 41st IEEE/ACM International Symposium on Microarchitecture Nov. 8, 2008:94-105. cited by applicant

Tiwari et al., Complete information flow tracking from the gates up. In Proceedings of the 14th international conference on Architectural support for programming languages and operating systems Mar. 7, 2009;109-20. cited by applicant

Waterman et al., The RISC-V Instruction Set Manual, vol. I: Unprivileged ISA, Document Version 20191213. RISC-V Foundation. Dec. 2019. 238 pages. cited by applicant

Watson et al., Cheri: A hybrid capability-system architecture for scalable software compartmentalization. 2015 IEEE Symposium on Security and Privacy Sep. 2015;20-37. cited by applicant

Witchel et al., Mondrian Memory Protection. ASPLOS-X. 2002; 13 pages. cited by applicant

Woodruff, Cheri: A RISC capability machine for practical memory safety. University of Cambridge, Computer Laboratory; 2014. 112 pages. cited by applicant

Zeldovich et al., Hardware Enforcement of Application Security Policies Using Tagged Memory. USENIX Symposium on Operating Systems Design and Implementation. Dec. 8, 2008:8;225-40. cited by applicant

U.S. Appl. No. 17/880,539, filed Aug. 3, 2022, Milburn et al. cited by applicant

U.S. Appl. No. 17/908,879, filed Sep. 1, 2022, Milburn et al. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2020/059057, mailed May 19, 2022. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/060030 mailed Mar. 19, 2020. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/060698 mailed Feb. 3, 2020. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/067084 mailed Mar. 19, 2020. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/013678 mailed Jun. 23, 2020. cited by applicant

International Preliminary Report on Patentability for International Application No. PCT/US2021/020602, mailed Sep. 15, 2022. cited by applicant

Primary Examiner: Henning; Matthew T

Attorney, Agent or Firm: Wolf, Greenfield & Sacks, P.C.

Background/Summary

RELATED APPLICATIONS (1) This application is a national stage filing under 35 U.S.C. § 371 of International Patent Application Serial No. PCT/US2020/013678, filed Jan. 15, 2020, which claims the benefit under 35 U.S.C. § 119(e) of U.S. Provisional Patent Application Ser. No. 62/794,499, filed on Jan. 18, 2019, titled “SYSTEMS AND METHODS FOR METADATA CLASSIFICATION”. The contents of these applications are incorporated herein by reference in their entirety.

BACKGROUND

(1) Computer security has become an increasingly urgent concern at all levels of society, from individuals to businesses to government institutions. For example, in 2015, security researchers identified a zero-day vulnerability that would have allowed an attacker to hack into a Jeep Cherokee's on-board computer system via the Internet and take control of the vehicle's dashboard functions, steering, brakes, and transmission. In 2017, the WannaCry ransomware attack was estimated to have affected more than 200,000 computers worldwide, causing at least hundreds of millions of dollars in economic losses. Notably, the attack crippled operations at several National Health Service hospitals in the UK. In the same year, a data breach at Equifax, a US consumer credit reporting agency, exposed person data such as full names, social security numbers, birth dates, addresses, driver's license numbers, credit card numbers, etc. That attack is reported to have affected over 140 million consumers.

(2) Security professionals are constantly playing catch-up with attackers. As soon as a vulnerability is reported, security professionals rush to patch the vulnerability. Individuals and organizations that fail to patch vulnerabilities in a timely manner (e.g., due to poor governance and/or lack of resources) become easy targets for attackers.

(3) Some security software monitors activities on a computer and/or within a network, and looks for patterns that may be indicative of an attack. Such an approach does not prevent malicious code from being executed in the first place. Often, the damage has been done by the time any suspicious pattern emerges.

SUMMARY

(4) In accordance with some embodiments, a method for metadata processing is provided, comprising acts of: receiving one or more metadata inputs; and processing the one or more metadata inputs to determine whether to allow an instruction, comprising: identifying one or more classification bits from a metadata input of the one or more metadata inputs; and processing the metadata input based on the one or more classification bits.

(5) In accordance with some embodiments, a method for metadata processing is provided, comprising acts of: receiving an application memory address; accessing mapping information associated with the application memory address; and using the mapping information to retrieve metadata from a memory region of a plurality of metadata memory regions, wherein the plurality of metadata memory regions correspond, respectively, to a plurality of metadata classifications.

- (6) In accordance with some embodiments, a method for metadata processing is provided, comprising acts of: identifying at least one metadata symbol from a metadata label; identifying at least one attribute value associated with the at least one metadata symbol; using the at least one attribute value to determine one or more classification bits; generating a binary representation for the metadata label, wherein: the binary representation comprises at least one portion representing the at least one metadata symbol; and the binary representation associates the one or more classification bits with the at least one portion representing the at least one metadata symbol.
- (7) In accordance with some embodiments, a method for metadata processing is provided, comprising acts of: identifying one or more registers and/or one or more application memory locations used by an instruction; accessing one or more classification bits for the one or more registers and/or one or more application memory locations; and determining, based on the one or more classification bits, whether to allow the instruction.
- (8) In accordance with some embodiments, a system is provided, comprising circuitry and/or one or more processors programmed by executable instructions, wherein the circuitry and/or the one or more programmed processors are configured to perform any of the methods described herein.
- (9) In accordance with some embodiments, at least one computer-readable medium is provided, having stored thereon at least one netlist for any of the circuitries described herein.
- (10) In accordance with some embodiments, at least one computer-readable medium is provided, having stored thereon at least one hardware description that, when synthesized, produces any of the netlists described herein.
- (11) In accordance with some embodiments, at least one computer-readable medium is provided, having stored thereon any of the executable instructions described herein.
-

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) FIG. 1 shows an illustrative hardware system **100** for enforcing policies, in accordance with some embodiments.
- (2) FIG. 2 shows an illustrative software system **200** for enforcing policies, in accordance with some embodiments.
- (3) FIG. 3 shows an illustrative register file **300**, in accordance with some embodiments.
- (4) FIG. 4 shows another example involving the illustrative register file **300** of FIG. 3, in accordance with some embodiments.
- (5) FIG. 5 shows an illustrative hardware accelerator **500**, in accordance with some embodiments.
- (6) FIG. 6 shows an illustrative hardware logic block **600**, in accordance with some embodiments.
- (7) FIG. 7A shows another example involving the illustrative register file **300** of FIG. 3, in accordance with some embodiments.
- (8) FIG. 7B shows another example involving the illustrative register file **300** of FIG. 3, in accordance with some embodiments.
- (9) FIG. 8 shows an illustrative hardware logic block **800**, in accordance with some embodiments.
- (10) FIG. 9 shows illustrative hardware logic blocks **900A-B**, in accordance with some embodiments.
- (11) FIG. 10 shows illustrative tag map table entries **1000A-B**, in accordance with some embodiments.
- (12) FIG. 11 shows an illustrative register file **1100**, in accordance with some embodiments.
- (13) FIG. 12 shows an illustrative hardware logic block **1200**, in accordance with some embodiments.
- (14) FIG. 13 shows, schematically, an illustrative computer **1000** on which any aspect of the present disclosure may be implemented.

DETAILED DESCRIPTION

(15) Many vulnerabilities exploited by attackers trace back to a computer architectural design where data and executable instructions are intermingled in a same memory. This intermingling allows an attacker to inject malicious code into a remote computer by disguising the malicious code as data. For instance, a program may allocate a buffer in a computer's memory to store data received via a network. If the program receives more data than the buffer can hold, but does not check the size of the received data prior to writing the data into the buffer, part of the received data would be written beyond the buffer's boundary, into adjacent memory. An attacker may exploit this behavior to inject malicious code into the adjacent memory. If the adjacent memory is allocated for executable code, the malicious code may eventually be executed by the computer.

(16) Techniques have been proposed to make computer hardware more security aware. For instance, memory locations may be associated with metadata for use in enforcing security policies, and instructions may be checked for compliance with the security policies. For example, given an instruction to be executed, metadata associated with the instruction and/or metadata associated with one or more operands of the instruction may be checked to determine if the instruction should be allowed. Additionally, or alternatively, appropriate metadata may be associated with an output of the instruction.

(17) FIG. 1 shows an illustrative hardware system **100** for enforcing policies, in accordance with some embodiments. In this example, the system **100** includes a host processor **110**, which may have any suitable instruction set architecture (ISA) such as a reduced instruction set computing (RISC) architecture or a complex instruction set computing (CISC) architecture. The host processor **110** may perform memory accesses via a write interlock **112**. The write interlock **112** may be connected to a system bus **115** configured to transfer data between various components such as the write interlock **112**, an application memory **120**, a metadata memory **125**, a read-only memory (ROM) **130**, one or more peripherals **135**, etc.

(18) In some embodiments, data that is manipulated (e.g., modified, consumed, and/or produced) by the host processor **110** may be stored in the application memory **120**. Such data is referred to herein as “application data,” as distinguished from metadata used for enforcing policies. The latter may be stored in the metadata memory **125**. It should be appreciated that application data may include data manipulated by an operating system (OS), instructions of the OS, data manipulated by one or more user applications, and/or instructions of the one or more user applications.

(19) In some embodiments, the application memory **120** and the metadata memory **125** may be physically separate, and the host processor **110** may have no access to the metadata memory **125**. In this manner, even if an attacker succeeds in injecting malicious code into the application memory **120** and causing the host processor **110** to execute the malicious code, the metadata memory **125** may not be affected. However, it should be appreciated that aspects of the present disclosure are not limited to storing application data and metadata on physically separate memories. Additionally, or alternatively, metadata may be stored in a same memory as application data, and a memory management component may be used that implements an appropriate protection scheme to prevent instructions executing on the host processor **110** from modifying the metadata. Additionally, or alternatively, metadata may be intermingled with application data in a same memory, and one or more policies may be used to protect the metadata.

(20) In some embodiments, tag processing hardware **140** may be provided to ensure that instructions being executed by the host processor **110** comply with one or more policies. The tag processing hardware **140** may include any suitable circuit component or combination of circuit components. For instance, the tag processing hardware **140** may include a tag map table **142** that maps addresses in the application memory **120** to addresses in the metadata memory **125**. For example, the tag map table **142** may map an address X in the application memory **120** to an address Y in the metadata memory **125**. A value stored at the address Y is sometimes referred to herein as a “metadata tag” or simply a “tag.”

(21) In some embodiments, a value stored at the address Y may in turn be an address Z. Such indirection may be repeated any suitable number of times, and may eventually lead to a data structure in the metadata memory **125** for storing metadata. Such metadata, as well as any intermediate address (e.g., the address Z), are also referred to herein as “metadata tags” or simply “tags.”

(22) It should be appreciated that aspects of the present disclosure are not limited to a tag map table that stores addresses in a metadata memory. In some embodiments, a tag map table entry itself may store metadata, so that the tag processing hardware **140** may be able to access the metadata without performing a memory operation. In some embodiments, a tag map table entry may store a selected bit pattern, where a first portion of the bit pattern may encode metadata, and a second portion of the bit pattern may encode an address in a metadata memory where further metadata may be stored. This may provide a desired balance between speed and expressivity. For instance, the tag processing hardware **140** may be able to check certain policies quickly, using only the metadata stored in the tag map table entry itself. For other policies with more complex rules, the tag processing hardware **140** may access the further metadata stored in the metadata memory **125**.

(23) Referring again to FIG. **1**, by mapping application memory addresses to metadata memory addresses, the tag map table **142** may create an association between application data and metadata that describes the application data. In one example, metadata stored at the metadata memory address Y and thus associated with application data stored at the application memory address X may indicate that the application data may be readable, writable, and/or executable. In another example, metadata stored at the metadata memory address Y and thus associated with application data stored at the application memory address X may indicate a type of the application data (e.g., integer, pointer, 16-bit word, 32-bit word, etc.). Depending on a policy to be enforced, any suitable metadata relevant for the policy may be associated with a piece of application data.

(24) In some embodiments, a metadata memory address Z may be stored at the metadata memory address Y. Metadata to be associated with the application data stored at the application memory address X may be stored at the metadata memory address Z, instead of (or in addition to) the metadata memory address Y. For instance, a binary representation of a metadata label RED may be stored at the metadata memory address Z. By storing the metadata memory address Z in the metadata memory address Y, the application data stored at the application memory address X may be tagged RED.

(25) In this manner, the binary representation of the metadata label RED may be stored only once in the metadata memory **125**. For instance, if application data stored at another application memory address X' is also to be tagged RED, the tag map table **142** may map the application memory address X' to a metadata memory address Y' where the metadata memory address Z is also stored.

(26) Moreover, in this manner, tag update may be simplified. For instance, if the application data stored at the application memory address X is to be tagged BLUE at a subsequent time, a metadata memory address Z' may be written at the metadata memory address Y, to replace the metadata memory address Z, and a binary representation of the metadata label BLUE may be stored at the metadata memory address Z'.

(27) Thus, the inventors have recognized and appreciated that a chain of metadata memory addresses of any suitable length N may be used for tagging, including N=0 (e.g., where a binary representation of a metadata label is stored at the metadata memory address Y itself).

(28) The association between application data and metadata (also referred to herein as “tagging”) may be done at any suitable level of granularity, and/or variable granularity. For instance, tagging may be done on a word-by-word basis. Additionally, or alternatively, a region in memory may be mapped to a single tag, so that all words in that region are associated with the same metadata. This may advantageously reduce a size of the tag map table **142** and/or the metadata memory **125**. For example, a single tag may be maintained for an entire address range, as opposed to maintaining multiple tags corresponding, respectively, to different addresses in the address range.

(29) In some embodiments, the tag processing hardware **140** may be configured to apply one or more rules to metadata associated with an instruction and/or metadata associated with one or more operands of the instruction to determine if the instruction should be allowed. For instance, the host processor **110** may fetch and execute an instruction, and may queue a result of executing the instruction into the write interlock **112**. Before the result is written back into the application memory **120**, the host processor **110** may send, to the tag processing hardware **140**, an instruction type (e.g., opcode), an address where the instruction is stored, one or more memory addresses referenced by the instruction, and/or one or more register identifiers. Such a register identifier may identify a register used by the host processor **110** in executing the instruction, such as a register for storing an operand or a result of the instruction.

(30) In some embodiments, destructive read instructions may be queued in addition to, or instead of, write instructions. For instance, subsequent instructions attempting to access a target address of a destructive read instruction may be queued in a memory region that is not cached. If and when it is determined that the destructive read instruction should be allowed, the queued instructions may be loaded for execution.

(31) In some embodiments, a destructive read instruction may be allowed to proceed, and data read from a target address may be captured in a buffer. If and when it is determined that the destructive read instruction should be allowed, the data captured in the buffer may be discarded. If and when it is determined that the destructive read instruction should not be allowed, the data captured in the buffer may be restored to the target address. Additionally, or alternatively, a subsequent read may be serviced by the buffered data.

(32) It should be appreciated that aspects of the present disclosure are not limited to performing metadata processing on instructions that have been executed by a host processor, such as instructions that have been retired by the host processor's execution pipeline. In some embodiments, metadata processing may be performed on instructions before, during, and/or after the host processor's execution pipeline.

(33) In some embodiments, given an address received from the host processor **110** (e.g., an address where an instruction is stored, or an address referenced by an instruction), the tag processing hardware **140** may use the tag map table **142** to identify a corresponding tag. Additionally, or alternatively, for a register identifier received from the host processor **110**, the tag processing hardware **140** may access a tag from a tag register file **146** within the tag processing hardware **140**.

(34) In some embodiments, if an application memory address does not have a corresponding entry in the tag map table **142**, the tag processing hardware **140** may send a query to a policy processor **150**. The query may include the application memory address in question, and the policy processor **150** may return a tag for that application memory address. Additionally, or alternatively, the policy processor **150** may create a new tag map entry for an address range including the application memory address. In this manner, the appropriate tag may be made available, for future reference, in the tag map table **142** in association with the application memory address in question.

(35) In some embodiments, the tag processing hardware **140** may send a query to the policy processor **150** to check if an instruction executed by the host processor **110** should be allowed. The query may include one or more inputs, such as an instruction type (e.g., opcode) of the instruction, a tag for a program counter, a tag for an application memory address from which the instruction is fetched (e.g., a word in memory to which the program counter points), a tag for a register in which an operand of the instruction is stored, and/or a tag for an application memory address referenced by the instruction. In one example, the instruction may be a load instruction, and an operand of the instruction may be an application memory address from which application data is to be loaded. The query may include, among other things, a tag for a register in which the application memory address is stored, as well as a tag for the application memory address itself. In another example, the instruction may be an arithmetic instruction, and there may be two operands. The query may include, among other things, a first tag for a first register in which a first operand is stored, and a

second tag for a second register in which a second operand is stored.

(36) It should also be appreciated that aspects of the present disclosure are not limited to performing metadata processing on a single instruction at a time. In some embodiments, multiple instructions in a host processor's ISA may be checked together as a bundle, for example, via a single query to the policy processor **150**. Such a query may include more inputs to allow the policy processor **150** to check all of the instructions in the bundle. Similarly, a CISC instruction, which may correspond semantically to multiple operations, may be checked via a single query to the policy processor **150**, where the query may include sufficient inputs to allow the policy processor **150** to check all of the constituent operations within the CISC instruction.

(37) In some embodiments, the policy processor **150** may include a configurable processing unit, such as a microprocessor, a field-programmable gate array (FPGA), and/or any other suitable circuitry. The policy processor **150** may have loaded therein one or more policies that describe allowed operations of the host processor **110**. In response to a query from the tag processing hardware **140**, the policy processor **150** may evaluate one or more of the policies to determine if an instruction in question should be allowed. For instance, the tag processing hardware **140** may send an interrupt signal to the policy processor **150**, along with one or more inputs relating to the instruction in question (e.g., as described above). The policy processor **150** may store the inputs of the query in a working memory (e.g., in one or more queues) for immediate or deferred processing. For example, the policy processor **150** may prioritize processing of queries in some suitable manner (e.g., based on a priority flag associated with each query).

(38) In some embodiments, the policy processor **150** may evaluate one or more policies on one or more inputs (e.g., one or more input tags) to determine if an instruction in question should be allowed. If the instruction is not to be allowed, the policy processor **150** may so notify the tag processing hardware **140**. If the instruction is to be allowed, the policy processor **150** may compute one or more outputs (e.g., one or more output tags) to be returned to the tag processing hardware **140**. As one example, the instruction may be a store instruction, and the policy processor **150** may compute an output tag for an application memory address to which application data is to be stored. As another example, the instruction may be an arithmetic instruction, and the policy processor **150** may compute an output tag for a register for storing a result of executing the arithmetic instruction.

(39) In some embodiments, the policy processor **150** may be programmed to perform one or more tasks in addition to, or instead of, those relating to evaluation of policies. For instance, the policy processor **150** may perform tasks relating to tag initialization, boot loading, application loading, memory management (e.g., garbage collection) for the metadata memory **125**, logging, debugging support, and/or interrupt processing. One or more of these tasks may be performed in the background (e.g., between servicing queries from the tag processing hardware **140**).

(40) In some embodiments, the tag processing hardware **140** may include a rule cache **144** for mapping one or more inputs to a decision and/or one or more outputs. For instance, a query into the rule cache **144** may be similarly constructed as a query to the policy processor **150** to check if an instruction executed by the host processor **110** should be allowed. If there is a cache hit, the rule cache **144** may output a decision as to whether the instruction should be allowed, and/or one or more output tags (e.g., as described above in connection with the policy processor **150**). Such a mapping in the rule cache **144** may be created using a query response from the policy processor **150**. However, that is not required, as in some embodiments, one or more mappings may be installed into the rule cache **144** ahead of time.

(41) In some embodiments, the rule cache **144** may be used to provide a performance enhancement. For instance, before querying the policy processor **150** with one or more input tags, the tag processing hardware **140** may first query the rule cache **144** with the one or more input tags. In case of a cache hit, the tag processing hardware **140** may proceed with a decision and/or one or more output tags from the rule cache **144**, without querying the policy processor **150**. This may provide a significant speedup. In case of a cache miss, the tag processing hardware **140** may query

the policy processor **150**, and may install a response from the policy processor **150** into the rule cache **144** for potential future use.

(42) In some embodiments, if the tag processing hardware **140** determines that an instruction in question should be allowed (e.g., based on a hit in the rule cache **144**, or a miss in the rule cache **144**, followed by a response from the policy processor **150** indicating no policy violation has been found), the tag processing hardware **140** may indicate to the write interlock **112** that a result of executing the instruction may be written back to memory. Additionally, or alternatively, the tag processing hardware **140** may update the metadata memory **125**, the tag map table **142**, and/or the tag register file **146** with one or more output tags (e.g., as received from the rule cache **144** or the policy processor **150**). As one example, for a store instruction, the metadata memory **125** may be updated based on an address translation by the tag map table **142**. For instance, an application memory address referenced by the store instruction may be used to look up a metadata memory address from the tag map table **142**, and metadata received from the rule cache **144** or the policy processor **150** may be stored to the metadata memory **125** at the metadata memory address. As another example, where metadata to be updated is stored in an entry in the tag map table **142** (as opposed to being stored in the metadata memory **125**), that entry in the tag map table **142** may be updated. As another example, for an arithmetic instruction, an entry in the tag register file **146** corresponding to a register used by the host processor **110** for storing a result of executing the arithmetic instruction may be updated with an appropriate tag.

(43) In some embodiments, if the tag processing hardware **140** determines that the instruction in question represents a policy violation (e.g., based on a miss in the rule cache **144**, followed by a response from the policy processor **150** indicating a policy violation has been found), the tag processing hardware **140** may indicate to the write interlock **112** that a result of executing the instruction should be discarded, instead of being written back to memory. Additionally, or alternatively, the tag processing hardware **140** may send an interrupt to the host processor **110**. In response to receiving the interrupt, the host processor **110** may switch to any suitable violation processing code. For example, the host processor **110** may halt, reset, log the violation and continue, perform an integrity check on application code and/or application data, notify an operator, etc.

(44) In some embodiments, the rule cache **144** may be implemented with a hash function and a designated portion of a memory (e.g., the metadata memory **125**). For instance, a hash function may be applied to one or more inputs to the rule cache **144** to generate an address in the metadata memory **125**. A rule cache entry corresponding to the one or more inputs may be stored to, and/or retrieved from, that address in the metadata memory **125**. Such an entry may include the one or more inputs and/or one or more corresponding outputs, which may be computed from the one or more inputs at run time, load time, link time, or compile time.

(45) In some embodiments, the tag processing hardware **140** may include one or more configuration registers. Such a register may be accessible (e.g., by the policy processor **150**) via a configuration interface of the tag processing hardware **140**. In some embodiments, the tag register file **146** may be implemented as configuration registers. Additionally, or alternatively, there may be one or more application configuration registers and/or one or more metadata configuration registers.

(46) Although details of implementation are shown in FIG. **1** and discussed above, it should be appreciated that aspects of the present disclosure are not limited to the use of any particular component, or combination of components, or to any particular arrangement of components. For instance, in some embodiments, one or more functionalities of the policy processor **150** may be performed by the host processor **110**. As an example, the host processor **110** may have different operating modes, such as a user mode for user applications and a privileged mode for an operating system. Policy-related code (e.g., tagging, evaluating policies, etc.) may run in the same privileged mode as the operating system, or a different privileged mode (e.g., with even more protection

against privilege escalation).

(47) FIG. 2 shows an illustrative software system **200** for enforcing policies, in accordance with some embodiments. For instance, the software system **200** may be programmed to generate executable code and/or load the executable code into the illustrative hardware system **100** shown in FIG. 1.

(48) In the example shown in FIG. 2, the software system **200** includes a software toolchain having a compiler **205**, a linker **210**, and a loader **215**. The compiler **205** may be programmed to process source code into executable code, where the source code may be in a higher-level language and the executable code may be in a lower level language. The linker **210** may be programmed to combine multiple object files generated by the compiler **205** into a single object file to be loaded by the loader **215** into memory (e.g., the illustrative application memory **120** in the example of FIG. 1). Although not shown, the object file output by the linker **210** may be converted into a suitable format and stored in persistent storage, such as flash memory, hard disk, read-only memory (ROM), etc. The loader **215** may retrieve the object file from the persistent storage, and load the object file into random-access memory (RAM).

(49) In some embodiments, the compiler **205** may be programmed to generate information for use in enforcing policies. For instance, as the compiler **205** translates source code into executable code, the compiler **205** may generate information regarding data types, program semantics and/or memory layout. As one example, the compiler **205** may be programmed to mark a boundary between one or more instructions of a function and one or more instructions that implement calling convention operations (e.g., passing one or more parameters from a caller function to a callee function, returning one or more values from the callee function to the caller function, storing a return address to indicate where execution is to resume in the caller function's code when the callee function returns control back to the caller function, etc.). Such boundaries may be used, for instance, during initialization to tag certain instructions as function prologue or function epilogue. At run time, a stack policy may be enforced so that, as function prologue instructions execute, certain locations in a call stack (e.g., where a return address is stored) may be tagged as “frame” locations, and as function epilogue instructions execute, the “frame” tags may be removed. The stack policy may indicate that instructions implementing a body of the function (as opposed to function prologue and function epilogue) only have read access to “frame” locations. This may prevent an attacker from overwriting a return address and thereby gaining control.

(50) As another example, the compiler **205** may be programmed to perform control flow analysis, for instance, to identify one or more control transfer points and respective destinations. Such information may be used in enforcing a control flow policy. As yet another example, the compiler **205** may be programmed to perform type analysis, for example, by applying type labels such as Pointer, Integer, Floating-Point Number, etc. Such information may be used to enforce a policy that prevents misuse (e.g., using a floating-point number as a pointer).

(51) Although not shown in FIG. 2, the software system **200** may, in some embodiments, include a binary analysis component programmed to take, as input, object code produced by the linker **210** (as opposed to source code), and perform one or more analyses similar to those performed by the compiler **205** (e.g., control flow analysis, type analysis, etc.).

(52) In the example of FIG. 2, the software system **200** further includes a policy compiler **220** and a policy linker **225**. The policy compiler **220** may be programmed to translate one or more policies written in a policy language into policy code. For instance, the policy compiler **220** may output policy code in C or some other suitable programming language. Additionally, or alternatively, the policy compiler **220** may output one or more metadata labels referenced by the one or more policies. At initialization, such a metadata label may be associated with one or more memory locations, registers, and/or other machine state of a target system, and may be resolved into a binary representation of metadata to be loaded into a metadata memory or some other hardware storage (e.g., registers) of the target system. As discussed above, such a binary representation of

metadata, or a pointer to a location at which the binary representation is stored, is sometimes referred to herein as a “tag.”

(53) It should be appreciated that aspects of the present disclosure are not limited to resolving metadata labels at load time. In some embodiments, one or more metadata labels may be resolved statically (e.g., at compile time or link time). For example, the policy compiler **220** may process one or more applicable policies, and resolve one or more metadata labels defined by the one or more policies into a statically determined binary representation. Additionally, or alternatively, the policy linker **225** may resolve one or more metadata labels into a statically determined binary representation, or a pointer to a data structure storing a statically determined binary representation. The inventors have recognized and appreciated that resolving metadata labels statically may advantageously reduce load time processing. However, aspects of the present disclosure are not limited to resolving metadata labels in any particular manner.

(54) In some embodiments, the policy linker **225** may be programmed to process object code (e.g., as output by the linker **210**), policy code (e.g., as output by the policy compiler **220**), and/or a target description, to output an initialization specification. The initialization specification may be used by the loader **215** to securely initialize a target system having one or more hardware components (e.g., the illustrative hardware system **100** shown in FIG. **1**) and/or one or more software components (e.g., an operating system, one or more user applications, etc.).

(55) In some embodiments, the target description may include descriptions of a plurality of named entities. A named entity may represent a component of a target system. As one example, a named entity may represent a hardware component, such as a configuration register, a program counter, a register file, a timer, a status flag, a memory transfer unit, an input/output device, etc. As another example, a named entity may represent a software component, such as a function, a module, a driver, a service routine, etc.

(56) In some embodiments, the policy linker **225** may be programmed to search the target description to identify one or more entities to which a policy pertains. For instance, the policy may map certain entity names to corresponding metadata symbols, and the policy linker **225** may search the target description to identify entities having those entity names. The policy linker **225** may identify descriptions of those entities from the target description, and use the descriptions to annotate, with appropriate metadata symbols, the object code output by the linker **210**. For instance, the policy linker **225** may apply a Read label to a .rodata section of an Executable and Linkable Format (ELF) file, a Read label and a Write label to a .data section of the ELF file, and an Execute label to a .text section of the ELF file. Such information may be used to enforce a policy for memory access control and/or executable code protection (e.g., by checking read, write, and/or execute privileges).

(57) It should be appreciated that aspects of the present disclosure are not limited to providing a target description to the policy linker **225**. In some embodiments, a target description may be provided to the policy compiler **220**, in addition to, or instead of, the policy linker **225**. The policy compiler **220** may check the target description for errors. For instance, if an entity referenced in a policy does not exist in the target description, an error may be flagged by the policy compiler **220**. Additionally, or alternatively, the policy compiler **220** may search the target description for entities that are relevant for one or more policies to be enforced, and may produce a filtered target description that includes entities descriptions for the relevant entities only. For instance, the policy compiler **220** may match an entity name in an “init” statement of a policy to be enforced to an entity description in the target description, and may remove from the target description (or simply ignore) entity descriptions with no corresponding “init” statement.

(58) In some embodiments, the loader **215** may initialize a target system based on an initialization specification produced by the policy linker **225**. For instance, referring to the example of FIG. **1**, the loader **215** may load data and/or instructions into the application memory **120**, and may use the initialization specification to identify metadata labels associated with the data and/or instructions

being loaded into the application memory **120**. The loader **215** may resolve the metadata labels in the initialization specification into respective binary representations. However, it should be appreciated that aspects of the present disclosure are not limited to resolving metadata labels at load time. In some embodiments, a universe of metadata labels may be known during policy linking, and therefore metadata labels may be resolved at that time, for example, by the policy linker **225**. This may advantageously reduce load time processing of the initialization specification. (59) In some embodiments, the policy linker **225** and/or the loader **215** may maintain a mapping of binary representations of metadata back to human readable versions of metadata labels. Such a mapping may be used, for example, by a debugger **230**. For instance, in some embodiments, the debugger **230** may be provided to display a human readable version of an initialization specification, which may list one or more entities and, for each entity, a set of one or more metadata labels associated with the entity. Additionally, or alternatively, the debugger **230** may be programmed to display assembly code annotated with metadata labels, such as assembly code generated by disassembling object code annotated with metadata labels. During debugging, the debugger **230** may halt a program during execution, and allow inspection of entities and/or metadata tags associated with the entities, in human readable form. For instance, the debugger **230** may allow inspection of entities involved in a policy violation and/or metadata tags that caused the policy violation. The debugger **230** may do so using the mapping of binary representations of metadata back to metadata labels.

(60) In some embodiments, a conventional debugging tool may be extended to allow review of issues related to policy enforcement, for example, as described above. Additionally, or alternatively, a stand-alone policy debugging tool may be provided.

(61) In some embodiments, the loader **215** may load the binary representations of the metadata labels into the metadata memory **125**, and may record the mapping between application memory addresses and metadata memory addresses in the tag map table **142**. For instance, the loader **215** may create an entry in the tag map table **142** that maps an application memory address where an instruction is stored in the application memory **120**, to a metadata memory address where metadata associated with the instruction is stored in the metadata memory **125**. Additionally, or alternatively, the loader **215** may store metadata in the tag map table **142** itself (as opposed to the metadata memory **125**), to allow access without performing any memory operation.

(62) In some embodiments, the loader **215** may initialize the tag register file **146** in addition to, or instead of, the tag map table **142**. For instance, the tag register file **146** may include a plurality of registers corresponding, respectively, to a plurality of entities. The loader **215** may identify, from the initialization specification, metadata associated with the entities, and store the metadata in the respective registers in the tag register file **146**.

(63) Referring again to the example of FIG. **1**, the loader **215** may, in some embodiments, load policy code (e.g., as output by the policy compiler **220**) into the metadata memory **125** for execution by the policy processor **150**. Additionally, or alternatively, a separate memory (not shown in FIG. **1**) may be provided for use by the policy processor **150**, and the loader **215** may load policy code and/or associated data into the separate memory.

(64) In some embodiments, a metadata label may be based on multiple metadata symbols. For instance, an entity may be subject to multiple policies, and may therefore be associated with different metadata symbols corresponding, respectively, to the different policies. The inventors have recognized and appreciated that it may be desirable that a same set of metadata symbols be resolved by the loader **215** to a same binary representation (which is sometimes referred to herein as a “canonical” representation). For instance, a metadata label {A, B, C} and a metadata label {B, A, C} may be resolved by the loader **215** to a same binary representation. In this manner, metadata labels that are syntactically different but semantically equivalent may have the same binary representation.

(65) The inventors have further recognized and appreciated it may be desirable to ensure that a

binary representation of metadata is not duplicated in metadata storage. For instance, as discussed above, the illustrative rule cache **144** in the example of FIG. **1** may map input tags to output tags, and, in some embodiments, the input tags may be metadata memory addresses where binary representations of metadata are stored, as opposed to the binary representations themselves. The inventors have recognized and appreciated that, if a same binary representation of metadata is stored at two different metadata memory addresses X and Y, the rule cache **144** may not recognize an input pattern having the metadata memory address Y as matching a stored mapping having the metadata memory address X. This may result in a large number of unnecessary rule cache misses, which may degrade system performance.

(66) Moreover, the inventors have recognized and appreciated that having a one-to-one correspondence between binary representations of metadata and their storage locations may facilitate metadata comparison. For instance, equality between two pieces of metadata may be determined simply by comparing metadata memory addresses, as opposed to comparing binary representations of metadata. This may result in significant performance improvement, especially where the binary representations are large (e.g., many metadata symbols packed into a single metadata label).

(67) Accordingly, in some embodiments, the loader **215** may, prior to storing a binary representation of metadata (e.g., into the metadata memory **125**), check if the binary representation of metadata has already been stored. If the binary representation of metadata has already been stored, instead of storing it again at a different storage location, the loader **215** may refer to the existing storage location. Such a check may be done at startup and/or when a program is loaded subsequent to startup (with or without dynamic linking).

(68) Additionally, or alternatively, a similar check may be performed when a binary representation of metadata is created as a result of evaluating one or more policies (e.g., by the illustrative policy processor **150**). If the binary representation of metadata has already been stored, a reference to the existing storage location may be used (e.g., installed in the illustrative rule cache **144**).

(69) In some embodiments, the loader **215** may create a hash table mapping hash values to storage locations. Before storing a binary representation of metadata, the loader **215** may use a hash function to reduce the binary representation of metadata into a hash value, and check if the hash table already contains an entry associated with the hash value. If so, the loader **215** may determine that the binary representation of metadata has already been stored, and may retrieve, from the entry, information relating to the binary representation of metadata (e.g., a pointer to the binary representation of metadata, or a pointer to that pointer). If the hash table does not already contain an entry associated with the hash value, the loader **215** may store the binary representation of metadata (e.g., to a register or a location in a metadata memory), create a new entry in the hash table in association with the hash value, and store appropriate information in the new entry (e.g., a register identifier, a pointer to the binary representation of metadata in the metadata memory, a pointer to that pointer, etc.). However, it should be appreciated that aspects of the present disclosure are not limited to the use of a hash table for keeping track of binary representations of metadata that have already been stored. Additionally, or alternatively, other data structures may be used, such as a graph data structure, an ordered list, an unordered list, etc. Any suitable data structure or combination of data structures may be selected based on any suitable criterion or combination of criteria, such as access time, memory usage, etc.

(70) It should be appreciated that the techniques introduced above and discussed in greater detail below may be implemented in any of numerous ways, as the techniques are not limited to any particular manner of implementation. Examples of implementation details are provided herein solely for illustrative purposes. Furthermore, the techniques disclosed herein may be used individually or in any suitable combination, as aspects of the present disclosure are not limited to the use of any particular technique or combination of techniques.

(71) For instance, while examples are discussed herein that include a compiler (e.g., the illustrative

compiler **205** and/or the illustrative policy compiler **220** in the example of FIG. 2), it should be appreciated that aspects of the present disclosure are not limited to using a compiler. In some embodiments, a software toolchain may be implemented as an interpreter. For example, a lazy initialization scheme may be implemented, where one or more default symbols (e.g., UNINITIALIZED) may be used for tagging at startup, and a policy processor (e.g., the illustrative policy processor **150** in the example of FIG. 1) may evaluate one or more policies and resolve the one or more default symbols in a just-in-time manner.

(72) The inventors have recognized and appreciated that certain patterns may arise in metadata processing. As an example, some policies (e.g., compartmentalization, heap protection, etc.) may use metadata symbols that are colloquially referred to as “colors.” Processing of such metadata symbols often involve comparisons (e.g., full or partial equality checking). For instance, a load/store instruction may be allowed to proceed if a color associated with an address register of the load/store instruction is the same as, or otherwise matches, a color associated with an application memory location referenced by an application memory address stored in the address register. Metadata involved in comparisons may be referred to herein as “comparison” metadata.

(73) As another example, metadata associated with an address (e.g., an application memory address, a register address, etc.) may sometimes flow with content stored at the address. Such metadata may be referred to herein as “flow” metadata. For instance, if the content is copied to a new address, the metadata may be carried over to the new address. Additionally, or alternatively, if the content is overwritten by new content from another address, the metadata may be overwritten by metadata associated with the other address.

(74) On the other hand, metadata associated with an address may sometimes stay with a location referenced by the address (e.g., one or more application memory cells, one or more registers, etc.). Such metadata may be referred to herein as “fixed” metadata. For instance, if content stored at the address is overwritten by new content, the metadata associated with the address may be preserved. Additionally, or alternatively, if the content is copied to a new address, the metadata may not be carried over.

(75) In some embodiments, flow metadata associated with an address may provide information about a data value stored at a location referenced by the address, whereas fixed metadata associated with the address may provide information about the location itself.

(76) The inventors have recognized and appreciated various advantages provided by classifying metadata based on processing patterns. For instance, in some embodiments, a binary representation of a metadata label may be augmented with one or more designated bits to indicate a metadata class to which the metadata label belongs (e.g., comparison, fixed, and/or flow). In this manner, one or more hardware components may be made aware of classification of metadata, which may allow certain processing logic to be implemented in hardware, thereby improving performance.

(77) FIG. 3 shows an illustrative register file **300**, in accordance with some embodiments. The register file **300** may be a register file of a host processor (e.g., the illustrative host processor **110** shown in FIG. 1), and may include a plurality of registers, such as registers **R0**, **R1**, etc.

(78) In some embodiments, the host processor **110** may execute a store instruction to write data to a location in an application memory (e.g., the illustrative application memory **120** shown in FIG. 1). An address (e.g., 0x1234) of the application memory location may be stored in an address register (e.g., **R0**), and the data to be written may be stored in a data register (e.g., **R1**).

(79) In some embodiments, each register in the register file **300** may have a corresponding tag register for storing metadata. The corresponding tag registers may be in a tag register file (e.g., the illustrative tag register file **146** shown in FIG. 1). For instance, the register **R0** may have a corresponding tag register in the tag register file **146**, and the corresponding tag register may store a binary representation <<RED>> of a metadata label RED. Likewise, the register **R1** may have a corresponding tag register in the tag register file **146**, and the corresponding tag register may store a binary representation <<BLUE>> of a metadata label BLUE.

(80) Additionally, or alternatively, the application memory address 0x1234 stored in the address register R0 may have a corresponding address in a metadata memory (e.g., the illustrative metadata memory 125 shown in FIG. 1). Metadata associated with the application memory address 0x1234 may be stored at the corresponding metadata memory address. For instance, in the example of FIG. 3, a binary representation <<GREEN>> of a metadata label GREEN may be stored at the metadata memory address corresponding to the application memory address 0x1234.

(81) In some embodiments, tag processing hardware (e.g., the illustrative tag processing hardware 140 shown in FIG. 1) may apply one or more policies to determine whether a store instruction should be allowed to write content of the data register R1 to the application memory address 0x1234. For instance, the tag processing hardware 140 may construct an input metadata pattern, and may use the input metadata pattern to query a policy processor and/or a rule cache (e.g., the illustrative policy processor 150 and/or the illustrative rule cache 144 shown in FIG. 1).

(82) For example, an access control policy may be provided that checks whether metadata associated with an address register matches metadata associated with an application memory location referenced by an application memory address stored in the address register. In the example of FIG. 3, the metadata label RED is associated with the address register R0, but the metadata label GREEN is associated with the application memory address 0x1234. Because of this mismatch, the tag processing hardware may indicate to the host processor that the store instruction violates the access control policy and therefore should not be allowed.

(83) In some embodiments, the rule cache 144 may store one or more entries for use in enforcing the illustrative access control policy described above. For instance, the following entries may be stored in the rule cache 144. 1) If R0=RED and mem=RED, then allow instruction 2) If R0=BLUE and mem=BLUE, then allow instruction 3) If R0=GREEN and mem=GREEN, then allow instruction

(84) The inventors have recognized and appreciated that, although the above access control policy involves only a simple equality check (namely, checking whether R0=mem), three entries may be stored in the rule cache because there are three distinct color values (namely, RED, GREEN, and BLUE).^{sup.1} More generally, if a policy uses N distinct color values, N entries may be stored in a rule cache to implement an equality check. ^{sup.1A} A rule cache entry (e.g., if R0=RED and mem=RED, then allow instruction) may be referred to herein as a “concrete” rule, whereas a policy rule (e.g., if R0=mem, then allow instruction) may be referred to herein as a “symbolic” rule.

(85) The inventors have further recognized and appreciated that, while a small number of color values (e.g., tens of color values) may be sufficient for some policies (e.g., compartmentalization), a large number of color values (e.g., tens of thousands of color values) may be used by other policies. For instance, for a heap protection policy, it may be desirable to provide temporal safety (e.g., preventing an attacker from re-using a stale pointer). Accordingly, a new color value may be added each time a portion of memory is dynamically allocated, and thus a potentially unbounded number of colors may be used.

(86) The inventors have recognized and appreciated that a bound may be imposed in practice based on a metadata tag size. For example, for 16-bit tags, ^{2.sup.16} 65,536 color values may be possible. However, it may be impractical to provide a rule cache that accommodates more than a few thousand entries, because there may be limited area available on a System-on-Chip (SoC). As a result, rule cache misses may occur frequently, and the policy processor 150 may be frequently invoked, which may negatively impact performance.

(87) In some embodiments, one or more policies (e.g., a heap protection policy) may be adapted to alias one or more color values. For instance, an equivalence relation may be provided over a set of color values, and a smaller set of color values may be obtained by taking a representative color value from each equivalence class. As an example, one or more test runs of relevant software may be conducted, and statistics may be collected regarding respective frequencies of color values. An equivalence relation may then be provided based on the observed frequencies. For instance, a

frequently occurring color value may be related only to itself, and is thus not aliased. By contrast, multiple color values that occur infrequently may be grouped together, and may be replaced by a representative color value from the group. In this manner, fewer color values may be used, and hence fewer rule cache misses may occur.

(88) The inventors have recognized and appreciated that, while aliasing color values of a policy may reduce a number of rule cache misses, an attacker may be able to exploit aliased color values to defeat a policy, and therefore there may be a negative impact on security. Accordingly, in some embodiments, one or more techniques are provided for improving performance without compromising security.

(89) In some embodiments, metadata may be classified based on how the metadata is created. For instance, a binary representation of a metadata label may be augmented with one or more designated bits to indicate whether the metadata label is dynamically created (e.g., as a result of enforcing one or more policies). Additionally, or alternatively, metadata may be classified based on a type of metadata processing that uses the metadata. For instance, a binary representation of a metadata label may be augmented with one or more designated bits to indicate whether a comparison (e.g., equal to, less than, less than or equal to, greater than, greater than or equal to, partially equal to, ones' complement of, etc.) is to be performed on the metadata label. Additionally, or alternatively, metadata may be classified based on how the metadata is to be processed. For instance, a binary representation of a metadata label may be augmented with one or more designated bits to indicate whether the metadata label is to be processed by a hardware accelerator (e.g., specialized circuitry adapted to check equality, inequality, partial equality, ones' complement, etc.). Any suitable combination of one or more classifications, or no classification at all, may be used, as aspects of the present disclosure are not so limited.

(90) FIG. 4 shows another example involving the illustrative register file **300** of FIG. 3, in accordance with some embodiments. In this example, the tag register corresponding to the address register **R0** stores a binary representation of the metadata label **RED**, augmented with a comparison bit that is set to 1. This is shown in FIG. 4 as <<RED>>*1, and may indicate that a comparison (e.g., equality checking) is to be performed on the metadata symbol **RED**. Likewise, the metadata memory address corresponding to the application memory address 0x1234 may store a binary representation of the metadata label **GREEN**, augmented with a comparison bit that is set to 1. This is shown in FIG. 4 as <<GREEN>>*1, and may indicate that a comparison (e.g., equality checking) is to be performed on the metadata symbol **GREEN**.

(91) It should be appreciated that aspects of the present disclosure are not limited to using comparison bits. Other classification bits, or no classification bit at all, may be used in addition to, or instead of, comparison bits. Moreover, despite the notation used herein (e.g., <<RED>>*1), aspects of the present disclosure are not limited to encoding classification bits in any particular manner. A classification bit may be appended, prepended, or otherwise encoded into a binary representation of metadata.

(92) FIG. 5 shows an illustrative hardware accelerator **500**, in accordance with some embodiments. For instance, the hardware accelerator **500** may include one or more hardware logic blocks inserted between a tag map table and a rule cache (e.g., the illustrative tag map table **142** and the illustrative rule cache **144** shown in FIG. 1). The one or more hardware logic blocks may be adapted to perform any suitable function, such as comparing two input metadata labels.

(93) Referring back to the example of FIG. 4, the tag map table **142** may map an address of the address register **R0** to an address of the corresponding tag register, and the metadata <<RED>>*1 may be read from the corresponding tag register. Additionally, or alternatively, the tag map table **142** may map the application memory address 0x1234 to the corresponding metadata memory address, and the metadata <<GREEN>>*1 may be read from the corresponding metadata memory address.

(94) In some embodiments, the metadata associated with the address register **R0** (e.g., <<RED>>*1

in the example of FIG. 4) and the metadata associated with the application memory address 0x1234 (e.g., <<GREEN>>*1 in the example of FIG. 4) may be provided to the hardware accelerator 500 as input. The hardware accelerator 500 may check if both metadata inputs include a comparison bit that is set to 1. If so, the hardware accelerator 500 may compare selected portions of the metadata inputs (e.g., <<RED>> and <<GREEN>>) to determine if there is a match, and may provide a comparison output accordingly. If at least one of the comparison bits is not set to 1, then the hardware accelerator 500 may forward the metadata inputs (or selected portions thereof, e.g., <<RED>> and <<GREEN>>, without comparison bits) to the rule cache 144 for further processing.

(95) In some embodiments, the hardware accelerator 500 may include specialized circuitry adapted to perform a certain type of comparison efficiently. For instance, the hardware accelerator 500 may include a logic block adapted to check if two input bit strings are the same. In this manner, equality checking may be implemented in hardware logic, instead of via the rule cache 144. For a policy that uses a large number of distinct color values (e.g., a heap protection policy with dynamically created color values), such hardware logic implementation may significantly reduce a number of rule cache misses and/or a physical size of the rule cache 144, and hence significantly improve performance and/or reduce chip area.

(96) It should be appreciated that any suitable set of metadata symbols may be involved in comparisons. Names of visually perceptible colors (e.g., RED, GREEN, BLUE, etc.) are used herein solely for purposes of illustration. Moreover, it should be appreciated that equality checking is discussed herein merely as an example. One or more other types of comparisons may be performed in addition to, or instead of, equality checking. For instance, in some embodiments, a comparison may be performed to determine if a bit string is a ones' complement of a second bit string. Additionally, or alternatively, a signed or unsigned arithmetic comparison (e.g., less than, less than or equal to, greater than, greater than or equal to, etc.) may be performed.

(97) It should also be appreciated that aspects of the present disclosure are not limited to having a hardware logic block adapted to perform comparisons. In some embodiments, the hardware accelerator 500 may include a plurality of hardware logic blocks adapted to perform different functions in addition to, or instead of, comparisons. These functional blocks may be activated independently, for example, based on one or more policies that are currently being enforced. For instance, each policy may have a corresponding functional block, which may be activated when the policy's code is loaded into a policy processor (e.g., the illustrative policy processor 150 shown in FIG. 1) for execution.

(98) FIG. 6 shows an illustrative hardware logic block 600, in accordance with some embodiments. For instance, the hardware logic block 600 may be part of the illustrative hardware accelerator 500 in the example of FIG. 5, and may be used to perform symbolic processing for a policy rule, such as a policy rule that receives three metadata labels as input, compares the first two input labels, and forwards the third input label to output if the first two input labels match. Thus, the hardware logic block 600 may provide a hardware implementation of a policy rule.

(99) In some embodiments, the hardware logic block 600 may receive one or more metadata inputs from tag fetch logic, which may include the tag map table 142. The one or more metadata inputs may be in binary form. In the example shown in FIG. 6, the hardware logic block 600 receives three inputs, L1, L2, and L3. For instance, the input L1 may be associated with an address register for a data transfer instruction, the input L2 may be associated with a source of the data transfer instruction (e.g., an application memory address for a load instruction, a data register for a store instruction, or a source address for a direct memory access instruction), and the input L3 may be associated with a target of the data transfer instruction (e.g., a data register for the load instruction, an application memory address for the store instruction, or a target address for the direct memory access instruction).

(100) Referring to the example of FIG. 4, the input L1 may be <<RED>>*1, the metadata read

from the tag register corresponding to the address register **R0**, whereas the input **L2** may be <<GREEN>>*1, the metadata read from the metadata memory address corresponding to the application memory address 0x1234, and the input **L3** may be <<BLUE>>*0, the metadata read from the tag register corresponding to the data register **R1**.

(101) In some embodiments, the hardware logic block **600** may include a sub-block **605**, which may be configured to determine whether symbolic processing is to be performed. For instance, the inputs **L1** and **L2** may each include a classification bit indicative of whether symbolic processing is to be performed, and the sub-block **605** may check if both classification bits have been set to 1. If both classification bits have been set to 1, a sub-block **610** may be activated to compare selected portions of the inputs **L1** and **L2** (e.g., <<RED>> and <<GREEN>>, without classification bits). Otherwise, the three inputs **L1**, **L2**, and **L3**, or selected portions thereof (e.g., <<RED>>, <<GREEN>>, and <<BLUE>>, without classification bits), may be forwarded to the rule cache **144**, which may check if there is an entry matching the three inputs. If there is no matching entry in the rule cache **144**, the policy processor **150** may be invoked on the inputs **L1**, **L2**, and **L3**, or the selected portions thereof. Additionally, or alternatively, the inputs **L1**, **L2**, and **L3**, or the selected portions thereof, may be forwarded to one or more accelerators other than the hardware accelerator **500**.

(102) In the example of FIG. 6, an output of the sub-block **605** is used to gate an input to the sub-block **610**, and an inverted output of the sub-block **605** is used to gate an input to the rule cache **144**. However, it should be appreciated that aspects of the present disclosure are not limited to any particular configuration of hardware logic blocks.

(103) In some embodiments, the sub-block **610** may be configured to perform an equality check on two input bit strings (e.g., <<RED>> and <<GREEN>>). If the input bit strings match, the input **L3**, a selected portion thereof (e.g., <<BLUE>>, without classification bits) may be forwarded along as an output of the hardware logic block **600**. For instance, an output of the sub-block **610** may be used to gate an output of the hardware logic block **600**. If the inputs **L1** and **L2** do not match, an error signal may be asserted, which may indicate that a policy has been violated. As a result, the store instruction attempting to write the data from the data register **R1** to the location referenced by the application memory address 0x1234 (which is stored in the address register **R0**) may be disallowed.

(104) In some embodiments, the hardware accelerator **500** may be configurable at run time to implement a plurality of different types of acceleration. For example, the sub-block **610** may be programmable, at run time, to switch among different types of matching, such as ==, >, <, >=, <=, etc. Additionally, or alternatively, an output sub-block may be provided to selectively forward one or more inputs (or selected portions thereof) in addition to, or instead of, **L3** (or selected portion thereof).

(105) In some embodiments, run time configurability may be provided via a lookup table, which may be implemented in any suitable manner (e.g., using a random-access memory). For instance, one or more values may be programmed into the lookup table at run time to indicate a desired acceleration operation corresponding to one or more classification bits. A key into the lookup table may be formed by concatenating, or otherwise combining, the one or more classification bits, and the lookup table may return the one or more values. However, it should be appreciated that aspects of the present disclosure are not limited to any particular technique for mapping a combination of one or more classification bits to a corresponding acceleration operation, or to using such a mapping at all. For instance, in some embodiments, an embedded FPGA fabric may be programmed to implement such a mapping.

(106) The inventors have recognized and appreciated that some policy rules have a simple logical structure (e.g., the illustrative logical structure shown in FIG. 6), and therefore may be implemented using a small amount of hardware logic. By processing such policy rules symbolically using hardware logic (instead of a rule cache), a number of rule cache misses and/or a

physical size of a rule cache may be reduced significantly, thereby significantly improving performance and/or reducing chip area.

(107) Although various applications and potential advantages are described above in connection with FIGS. 5-6, it should be appreciated that aspects of the present disclosure are not limited to using a hardware accelerator. In some embodiments, one or more of the hardware logic blocks of the hardware accelerator **500** may be included in the tag map table **142** and/or the rule cache **144**. Moreover, aspects of the present disclosure are not limited to using a classification bit to choose between a hardware accelerator and a rule cache. Additionally, or alternatively, one or more classification bits may be used to determine how a rule cache lookup is to be performed (e.g., which inputs are to be presented, to which rule cache the inputs are to be presented in case multiple rule caches are used, etc.).

(108) It should also be appreciated that aspects of the present disclosure are not limited to processing metadata associated with data transfer instructions. For instance, in some embodiments, the inputs **L1**, **L2**, and **L3** may be associated, respectively, with registers holding operands of an arithmetic instruction.

(109) FIG. 7A shows another example involving the illustrative register file **300** of FIG. 3, in accordance with some embodiments. In this example, the metadata label associated with the application memory address 0x1234 is RED, instead of GREEN. Thus, the illustrative access control rule described above may allow the store instruction to proceed, so that the content of the data register **R1** may be written to the application memory address 0x1234.

(110) In some embodiments, an information flow rule may be provided to propagate metadata associated with a data register to an application memory address to which data from the data register is written. In the example of FIG. 7A, the metadata label associated with the data register **R1** (namely, BLUE) may, as a result of the store instruction, be carried over to the metadata memory address corresponding to the application memory address 0x1234, replacing the metadata label RED.

(111) This may, however, interfere with the illustrative access control rule described above. For instance, associating the application memory address 0x1234 with the metadata label RED may mean that the application memory location referenced by the address 0x1234 is allocated for an application process A, and therefore the metadata label RED should remain with the address 0x1234 until the application memory location is released by the application process A, regardless of what content is stored in the application memory location. However, the illustrative information flow rule described above may replace the metadata label RED with the metadata label associated with the data register **R1** (namely, BLUE), even though the application memory location referenced by the address 0x1234 has not been released by the application process A.

(112) In some embodiments, interference between different rules may be avoided by using separate sets of metadata symbols. For instance, a policy language may be provided that allows metadata symbols to be declared using algebraic data types. As an example, there may be a metadata type Color, which may have three values: RED, BLUE, and GREEN. A constructor PTR may be applied to the metadata type Color to obtain a new metadata type PTR Color, which may also have three values: PTR RED, PTR BLUE, and PTR GREEN. Intuitively, an association with such a metadata symbol may be interpreted as a statement, “I am a red pointer,” “I am a blue pointer,” or “I am a green pointer.”

(113) Additionally, or alternatively, a constructor CEL may be applied to the metadata type Color to obtain a new metadata type CEL Color, which may also have three values: CEL RED, CEL BLUE, and CEL GREEN. Intuitively, an association with such a metadata symbol may be interpreted as a statement, “I reside at a red memory cell,” “I reside at a blue memory cell,” or “I reside at a green memory cell.”

(114) FIG. 7B shows another example involving the illustrative register file **300** of FIG. 3, in accordance with some embodiments. In this example, separate sets of metadata symbols are used.

For instance, instead of the metadata symbol RED, the address register R0 may be associated with the metadata symbol PTR RED, indicating that the application memory address 0x1234 stored in the address register R0 is a red pointer. Similarly, instead of the metadata symbol BLUE, the data register R1 may be associated with the metadata symbol PTR BLUE, indicating that the content of the data register R1 (which may be another application memory address, namely, 0x5678) is a blue pointer. Moreover, the application memory address 0x1234 may be associated with the metadata symbols CEL RED and PTR GREEN, indicating that the application memory address 0x1234 points to a red memory cell, and that a green pointer is stored therein.

(115) In some embodiments, when a store instruction attempts to write the content of the data register R1 to the location referenced by the application memory address 0x1234, an access control rule may allow the store instruction to proceed, because the metadata associated with the address register R0 (namely, PTR RED) matches the metadata associated with the application memory address 0x1234 (namely, CEL RED).

(116) On the other hand, an information flow rule may propagate the metadata associated with the data register R1 (namely, PTR BLUE) to the metadata memory address corresponding to the application memory address 0x1234, replacing the metadata symbol PTR GREEN. As a result, upon execution of the store instruction, the application memory address 0x1234 may be associated with the metadata symbol PTR BLUE, without disturbing the metadata symbol CEL RED. In this manner, the information flow rule may not interfere with the access control rule.

(117) The inventors have recognized and appreciated that the metadata symbols PTR RED (“I am a red pointer”), PTR BLUE (“I am a blue pointer”), and PTR GREEN (“I am a green pointer”) in the example of FIG. 7B may provide information about content that is stored at a location, whereas the metadata symbols CEL RED (“I reside at a red memory cell”), CEL BLUE (“I reside at a blue memory cell”), and CEL GREEN (“I reside at a blue memory cell”) may provide information about the location itself.

(118) The inventors have further recognized and appreciated that metadata that provides information about content and metadata that provides information about a location may have different update semantics. For example, in some instances, a piece of metadata may indicate that current content of an application memory location is considered private data of a particular user or application instance. Such content may be treated as private data regardless of how the content subsequently travels, and therefore the metadata may “flow” with the content. For instance, if the content is copied to a new location, the metadata may be carried over to an address of the new location. Additionally, or alternatively, if the content is overwritten by new content from another location, the metadata may be overwritten by metadata associated with the new content (e.g., metadata indicating that the new content is considered public data).

(119) By contrast, in other instances, a piece of metadata may indicate that an application memory location is allocated for use by a particular user or application instance. Such metadata may stay “fixed” at the application memory location, regardless of what content is stored there. For instance, the metadata may be preserved even if content stored at the location is overwritten by new content. Additionally, or alternatively, if the content is copied to a new location, the metadata may not be carried over.

(120) The inventors have recognized and appreciated that it may be desirable to make one or more hardware components aware of whether a metadata label includes flow metadata, fixed metadata, or both. This may allow a hardware component to process the metadata label differently depending on how the metadata label is classified, which may in turn lead to improved performance.

(121) However, in practice, it may be challenging to implement this type of differentiated processing via hardware logic. For instance, metadata is often stored in an encoded form. It may be inefficient to decode a metadata label to determine whether the metadata label includes flow metadata, fixed metadata, or both (e.g., private data stored at a private location). Moreover, although the constructors PTR and CEL may be used to differentiate between flow metadata and

fixed metadata for the illustrative heap protection policy discussed above, in practice a large number of policies may be enforced at the same time, and some of the policies may be loaded dynamically. Therefore, it may be challenging to keep track of such distinctions for all policies at a hardware level.

(122) Accordingly, in some embodiments, a policy language may be provided that allows attribute values to be declared for metadata symbols. For instance, a metadata symbol such as PTR RED may be declared with an attribute value that indicates the metadata symbol is flow metadata, whereas a metadata symbol such as CEL RED may be declared with an attribute value that indicates the metadata symbol is fixed metadata.

(123) In some embodiments, an attribute value may be declared for a metadata type. For example, with reference to the illustrative algebraic data types discussed above in connection with FIG. 7A, an attribute value may be declared for the metadata type Color to indicate that a metadata symbol of the type Color (e.g., RED, BLUE, GREEN, etc.) is comparison metadata. Additionally, or alternatively, an attribute value may be declared for the constructor PTR to indicate that a metadata symbol of the type PTR Color (e.g., PTR RED, PTR BLUE, PTR GREEN, etc.) is flow metadata. Additionally, or alternatively, an attribute value may be declared for the constructor CEL to indicate that a metadata symbol of the type CEL Color (e.g., CEL RED, CEL BLUE, CEL GREEN, etc.) is fixed metadata.

(124) In some embodiments, a policy compiler and/or a policy linker (e.g., the illustrative policy compiler **220** and/or the illustrative policy linker **225** in the example of FIG. 2) may assemble attribute values for a composite metadata type. For instance, the policy compiler **220** and/or the policy linker **225** may mark the metadata symbol CEL RED with both the attribute value indicative of fixed metadata and the attribute value indicative of comparison metadata, because the constructor CEL is marked with the former, and the metadata symbol RED is marked with the latter. Likewise, the policy compiler **220** and/or the policy linker **225** may mark the metadata symbol PTR RED with both the attribute value indicative of flow metadata and the attribute value indicative of comparison metadata, because the constructor PTR is marked with the former, and the metadata symbol RED is marked with the latter.

(125) Additionally, or alternatively, the policy compiler **220** and/or the policy linker **225** may assemble attribute values for a composite metadata label, which may include a set of metadata symbols for respective policies that are being enforced simultaneously. For instance, the policy compiler **220** and/or the policy linker **225** may mark the composite metadata label with a set of attribute values, where each attribute value in the set is used to mark some metadata symbol in the composite metadata label. With reference to the example of FIG. 7B, the application memory address 0x1234 may be associated with a composite metadata label {PTR GREEN, CEL RED}, where the metadata symbol PTR GREEN may be used by an information flow rule, whereas the metadata symbol CEL RED may be used by an access control rule. The policy compiler **220** and/or the policy linker **225** may mark this composite metadata label with all three attribute values: flow and comparison from PTR GREEN, and fixed and comparison from CEL RED.

(126) It should be appreciated that access control rules and information flow rules are described herein solely for purposes of illustration. Flow metadata may be used by any suitable rule to determine how to update metadata, and fixed metadata may be used by any suitable rule to determine whether to allow an instruction.

(127) In some embodiments, the policy compiler **220** and/or the policy linker **225** may use one or more attribute values for a metadata label to determine how to construct a binary representation of metadata. For instance, a binary representation of the metadata label may be augmented with one or more bits that encode one or more respective attribute values, such as a first bit for flow, a second bit for fixed, a third bit for comparison, etc. Thus, a metadata label such as PTR RED may have a binary representation with designated bits 101 (for flow and comparison), which may be denoted <<PTR RED>>*101. Likewise, a metadata label such as CEL RED may have a binary

representation with designated bits 011 (for fixed and comparison), which may be denoted <<CEL RED>>*011. A composite metadata label such as {PTR GREEN, CEL RED} may have a binary representation with designated bits 111 (for flow, fixed, and comparison), which may be denoted <<{PTR GREEN, CEL RED}>>*111. Additionally, or alternatively, respective designated bits may be maintained for component metadata symbols, such as {<<PTR GREEN>>*101, <<CEL RED>>*011}.

(128) The inventors have recognized and appreciated that such designated bits may be readily decoded by hardware logic to allow differentiated processing, thereby improving performance. However, it should be appreciated that aspects of the present disclosure are not limited to using any designated bit to perform differentiated processing at a hardware level, or to any particular combination of one or more designated bits. For instance, returning to the example of FIG. 7B, two designated bits are shown: 00 for none, 01 for fixed, 10 for flow, and 11 for both.

(129) FIG. 8 shows an illustrative hardware logic block **800**, in accordance with some embodiments. For instance, the hardware logic block **800** may be used to provide differentiated processing of metadata based on one or more classification bits. In some embodiments, the hardware logic block **800** may be provided as part of the illustrative hardware accelerator **500** shown in FIG. 5. However, that is not required. In some embodiments, one or more functionalities of the hardware logic block **800** may instead be implemented by a tag map table and/or a rule cache (e.g., the illustrative tag map table **142** and/or the illustrative rule cache **144** shown in FIG. 1).

(130) In the example of FIG. 8, the hardware logic block **800** receives three inputs, L1, L2, and L3, from tag fetch logic, which may include the tag map table **142**. These inputs may be in binary form. In some instances, the input L1 may be associated with an address register for a data transfer instruction, the input L2 may be associated with a source of the data transfer instruction (e.g., an application memory address for a load instruction, a data register for a store instruction, or a source address for a direct memory access instruction), and the input L3 may be associated with a target of the data transfer instruction (e.g., a data register for the load instruction, an application memory address for the store instruction, or a target address for the direct memory access instruction).

(131) Referring to the example of FIG. 7B, the input L1 may be <<PTR RED>>*10, the metadata read from the tag register corresponding to the address register R0, whereas the input L2 may be <<PTR BLUE>>*10, the metadata read from the tag register corresponding to the data register R1, and the input L3 may be {<<PTR GREEN>>*10, <<CEL RED>>*01}, the metadata read from the metadata memory address corresponding to the application memory address 0x1234.

(132) In some embodiments, the hardware logic block **800** may include a sub-block **805**, which may be configured to check if the inputs L1, L2, and L3 all include designated bits 00 (neither flow nor fixed). This may be performed expeditiously, without fully decoding L1, L2, and L3. If L1, L2, and L3 all include designated bits 00, no further metadata processing may be performed, and the data transfer instruction may be allowed.

(133) The inventors have recognized and appreciated that, in many instances, a host processor (e.g., the illustrative host processor **110** shown in FIG. 1) may be processing application data with neither flow metadata nor fixed metadata. For instance, when application data that is associated with fixed metadata, but no flow metadata, is loaded into a data register, the fixed metadata is not carried over, and therefore the application data, residing in the data register, may have neither flow metadata nor fixed metadata. Subsequently, the host processor **110** may perform arithmetic instructions on the application data with neither flow metadata nor fixed metadata. The inventors have recognized and appreciated that the sub-block **805** may allow such instructions to be quickly identified and disposed of without further metadata processing (e.g., without any rule cache lookup), which may in turn allow tag processing hardware (e.g., the illustrative tag processing hardware **140** shown in FIG. 1) to catch up following operations that take more than one clock cycle.

(134) For example, in some instances, the tag processing hardware **140** may have yet to retrieve metadata (e.g., from the illustrative metadata memory **125** shown in FIG. 1) for a first instruction

executed by the host processor **110**, even though the host processor **110** has already executed at least one second instruction after the first instruction. Disposing of the at least one second instruction expeditiously, based on designated bits, may keep the tag processing hardware **140** from falling too far behind the host processor **110**.

(135) In some embodiments, if not all of **L1**, **L2**, and **L3** include designated bits 00, a sub-block **810** may be activated to check if any of the inputs **L1**, **L2**, or **L3** includes designated bits 10 (flow). For instance, an inverted output of the sub-block **805** may be used to gate an input to the sub-block **810**.

(136) The inventors have recognized and appreciated that, if none of **L1**, **L2**, and **L3** includes flow metadata, there may be no need to perform any metadata update. For instance, referring to the example of FIG. 7A, if no information flow rule is being enforced, the metadata label BLUE associated with the data register **R1** may not be marked with designated bits 10 (flow), and may not be carried over as a result of the store instruction. Therefore, the metadata associated with the application memory address 0x1234 may not be updated following the store instruction (e.g., the metadata label RED may remain).

(137) As another example, a heap protection policy may be enforced, but neither the data register **R1** nor the memory location at the application memory address 0x1234 may store a pointer. Thus, neither the data register **R1** nor the application memory address 0x1234 may be associated with any flow metadata of the heap protection policy, and no metadata update may be performed as a result of the store instruction.

(138) Accordingly, in some embodiments, the sub-block **810** may check, expeditiously, if any of the inputs **L1**, **L2**, or **L3** includes designated bits 10 (flow), without fully decoding **L1**, **L2**, and **L3**. If none of **L1**, **L2**, and **L3** is marked with 10 (flow), **L1**, **L2**, and **L3** (or selected portions thereof) may be forwarded to a fixed rule cache **144A**, which may be configured to process fixed metadata only. For instance, an inverted output of the sub-block **810** may be used to gate an input to the fixed rule cache **144A**.

(139) The inventors have recognized and appreciated that, because there is no metadata update, the fixed rule cache **144A** may not output any metadata label, and therefore may be narrower than the rule cache **144**. Thus, by using a separate rule cache configured to process fixed metadata only, memory savings may be achieved. However, it should be appreciated that aspects of the present disclosure are not limited to using a separate rule cache.

(140) In some embodiments, if at least one of **L1**, **L2**, and **L3** includes designated bits 10 (flow), a sub-block **815** may be activated. For instance, an output of the sub-block **810** may be used to gate an input to the sub-block **815**.

(141) The inventors have recognized and appreciated that, if the input **L3**, which may be associated with the target of the data transfer instruction, does not include fixed metadata, there may be no need to perform any rule cache lookup. For instance, referring again to the example of FIG. 7A, if no access control rule is being enforced, the metadata label RED associated with the application memory address 0x1234 may not be marked with designated bits 01 (fixed), and the store instruction may be allowed without performing any rule cache lookup. As a result of the store instruction, the metadata label BLUE associated with the data register **R1** may be carried over, without being fully decoded, to the metadata memory address associated with the application memory address 0x1234. This may be appropriate because an information flow rule may simply indicate that flow metadata is to be propagated from source to target. Precise meaning of the flow metadata may be irrelevant for such an information flow rule.

(142) Accordingly, in some embodiments, the sub-block **815** may check, expeditiously, if **L3** includes designated bits 01 (fixed), without fully decoding **L3**. If **L3** does not include designated bits 01 (fixed), the sub-block **815** may check if **L1** and **L2** are reducible. For instance, the sub-block **815** may check if only one of **L1** and **L2** includes designated bits 10 (flow). Additionally, or alternatively, the sub-block **815** may check if both **L1** and **L2** include designated bits 10 (flow) but

include the same flow metadata.

(143) If L3 does not include designated bits 01 (fixed) and only one of L1 and L2 includes designated bits 10 (flow), then that one of L1 and L2 (or a selected portion thereof) may be forwarded along as an output of the hardware logic block **800**. Additionally, or alternatively, if L3 does not include designated bits 01 (fixed) and both L1 and L2 include designated bits 10 (flow) but include the same flow metadata, then either L1 or L2 (or a selected portion thereof) may be forwarded along as an output of the hardware logic block **800**, provided L3 does not include designated bits 01 (fixed).

(144) In some embodiments, L1 or L2 (or a selected portion thereof) may be forwarded along as an output of the hardware logic block **800** without being fully decoded. Again, this may be appropriate because an information flow rule may simply indicate that flow metadata is to be propagated from source to target, without regard to precise meaning of the flow metadata. For instance, the input L1 may be associated with a register storing a first operand of an arithmetic instruction (e.g., adding a constant offset to a memory address), the input L2 may be associated with a register storing a second operand of the arithmetic instruction, and the input L3 may be associated with a register storing a result of the arithmetic instruction. After execution of the instruction, flow metadata may simply be propagated from one of the operands (e.g., the memory address) to the result.

(145) In some embodiments, an output of the sub-block **815** may be used to gate an output of the hardware logic block **800**. For instance, if L3 does not include designated bits 01 (fixed), the sub-block **815** may output a signal to enable the output of the hardware logic block **800**. Additionally, or alternatively, the sub-block **815** may output a signal to select between L1 and L2 as the output of the hardware logic block **800**.

(146) In some embodiments, if one or more checks performed by the sub-block **815** fail, L1, L2, and L3 (or selected portions thereof) may be forwarded to the rule cache **144** for further processing. For instance, an inverted output of the sub-block **815** may be used to gate an input to the rule cache **144**.

(147) It should be appreciated that details of implementation are shown in FIG. **8** and described above solely for purposes of illustration. For instance, the hardware logic block **800** may include any suitable combination of one or more of the sub-blocks **805**, **810**, and **815**, and/or one or more other sub-blocks. Moreover, aspects of the present disclosure are not limited to any particular ordering of the sub-blocks **805**, **810**, and **815**. In some embodiments, the sub-blocks **805**, **810**, and **815** may be placed in parallel, instead of being in a sequence.

(148) It should also be appreciated that the techniques described herein may be implemented in any suitable manner, for example, using any suitable combination of software and/or hardware. For instance, any one or more of the checks performed by the hardware logic block **800** may be implemented via software executed by a policy processor (e.g., the illustrative policy processor **150** shown in FIG. **1**).

(149) In some embodiments, a single rule may incorporate both access control and information flow aspects. Referring to the example of FIG. **7B**, a combined rule may check whether the content metadata associated with the address register R0 (namely, PTR RED) matches the location metadata associated with the application memory address 0x1234 (namely, CEL RED). If so, the content metadata associated with the data register R1 (namely, PTR BLUE) may be propagated to the metadata memory address associated with the application memory address 0x1234.

(150) The inventors have recognized and appreciated that, while using a combined rule may result in fewer invocations of a policy processor and/or fewer accesses to a rule cache, there may be higher rule cache pressure. For instance, in the example of FIG. **7B**, with a combined rule and three colors RED, BLUE, and GREEN, 3.sup.2=9 rule cache entries may be stored as follows. 1) If R0=PTR RED and mem={CEL RED, DEF} and R1=PTR RED then mem={CEL RED, PTR RED} 2) If R0=PTR RED and mem={CEL RED, DEF} and R1=PTR BLUE then mem={CEL RED, PTR BLUE} 3) If R0=PTR RED and mem={CEL RED, DEF} and R1=PTR GREEN then

mem={CEL RED, PTR GREEN} 4) If R0=PTR BLUE and mem={CEL BLUE, DEF} and R1=PTR RED then mem={CEL BLUE, PTR RED} 5) If R0=PTR BLUE and mem={CEL BLUE, DEF} and R1=PTR BLUE then mem={CEL BLUE, PTR BLUE} 6) If R0=PTR BLUE and mem={CEL BLUE, DEF} and R1=PTR GREEN then mem={CEL BLUE, PTR GREEN} 7) If R0=PTR GREEN and mem={CEL GREEN, DEF} and R1=PTR RED then mem={CEL GREEN, PTR RED} 8) If R0=PTR GREEN and mem={CEL GREEN, DEF} and R1=PTR BLUE then mem={CEL GREEN, PTR BLUE} 9) If R0=PTR GREEN and mem={CEL GREEN, DEF} and R1=PTR GREEN then mem={CEL GREEN, PTR GREEN}

(151) In this example, the symbol DEF denotes a default metadata value, such as a bit string of all 1's. For instance, the content metadata associated with the application memory address 0x1234 prior to the store instruction (namely, PTR GREEN) may be irrelevant in deciding whether the store instruction should be allowed, and may be replaced by the content metadata associated with the data register R1 (namely, PTR BLUE) as a result of the store instruction. Therefore, the content metadata associated with the application memory address 0x1234 prior to the store instruction may be masked (e.g., replaced by the default symbol DEF) before being presented to a rule cache (e.g., the illustrative rule cache 144 in the example of FIG. 1).

(152) The inventors have recognized and appreciated that the combined rule described above may lead to increased rule cache pressure because the content metadata associated with the address register R0 and the content metadata associated with the data register R1 may be independent of each other. For instance, the content metadata associated with the address register R0 and the content metadata associated with the data register R1 may independently take on one of three distinct color values (namely, PTR RED, PTR BLUE, and PTR GREEN), so that a total of $3 \times 3 = 9$ rule cache entries may arise. More generally, if a rule has M independent variables, where each variable may take on one of N distinct values, a total of N^M entries may be stored in the rule cache.

(153) By contrast, with M separate rules each involving just one variable, $M \times N$ rule cache entries may be sufficient. Referring again to the example of FIG. 7B, with separate rules for access control and information flow, $2 \times 3 = 6$ rule cache entries may be provided as follows. 1) If R0=PTR RED and mem=CEL RED, then allow store instruction 2) If R0=PTR BLUE and mem=CEL BLUE, then allow store instruction 3) If R0=PTR GREEN and mem=CEL GREEN, then allow store instruction 1) If R1=PTR RED, then mem=PTR RED 2) If R1=PTR BLUE, then mem=PTR BLUE 3) If R1=PTR GREEN, then mem=PTR GREEN

(154) The inventors have recognized and appreciated that, where M is constant (e.g., two variables, one for access control and the other for information flow), $M \times N$ grows more slowly than N^M as N (the number of distinct color values) grows. Therefore, it may be possible to reduce rule cache pressure by breaking down a rule involving M independent variables into multiple rules that each involve a smaller number of independent variables.

(155) However, the inventors have further recognized and appreciated that, if separate access control and information flow rules are enforced via composite rule cache entries, a reduction in cache pressure may not be achieved. For instance, a composite rule cache entry may be constructed as a matrix of M rows, where each row corresponds to a respective policy rule and stores a constituent rule cache entry for that policy rule. Referring to the illustrative access control and information flow rules described above, there may be M=2 rows in each composite rule cache entry, one for the access control rule and one for the information flow rule. Since each row has N=3 possibilities, there may be $N^M = 3^2 = 9$ possible composite rule cache entries. That is the same number of rule cache entries for the illustrative combined rule described above, so there is no reduction.

(156) Accordingly, in some embodiments, separate policy rules (e.g., an access control rule and an information flow rule) may be enforced using separate rule cache entries. Techniques are provided herein for managing multiple rule cache lookups, so that an instruction may be allowed only if all

policy rules represented by the separate rule cache entries are satisfied. For instance, in some embodiments, one or more hardware logic blocks may be provided to use one or more metadata labels to construct inputs for the different rule cache lookups.

(157) FIG. 9 shows illustrative hardware logic blocks **900A-B**, in accordance with some embodiments. For instance, the hardware logic blocks **900A-B** may be used to construct inputs for different rule cache lookups. In some embodiments, the hardware logic blocks **900A-B** may be provided as part of the illustrative hardware accelerator **500** shown in FIG. 5. However, that is not required. In some embodiments, one or more functionalities of the hardware logic blocks **900A-B** may instead be implemented by a tag map table and/or a rule cache (e.g., the illustrative tag map table **142** and/or the illustrative rule cache **144** shown in FIG. 1).

(158) Similar to the illustrative hardware logic block **800** in the example of FIG. 8, the hardware logic blocks **900A-B** may each receive three metadata inputs from tag fetch logic, which may include the tag map table **142**. Referring to the example of FIG. 7B, the first input may be <<PTR RED>>*10, the metadata read from the tag register corresponding to the address register **R0**, whereas the second input may be <<PTR BLUE>>*10, the metadata read from the tag register corresponding to the data register **R1**, and the third input may be {<<PTR GREEN>>*10, <<CEL RED>>*01}, the metadata read from the metadata memory address corresponding to the application memory address 0x1234.

(159) In some embodiments, binary representations of metadata may include designated bits indicative of how metadata symbols are classified, such as 00 for none, 01 for fixed, 10 for flow, and 11 for both. For instance, the metadata symbol PTR RED associated with the address register **R0**, the metadata symbol PTR BLUE associated with the data register **R1**, and the metadata symbol PTR GREEN associated with the application memory address 0x1234 may be marked with 10 (flow), whereas the metadata symbol CEL RED associated with the application memory address 0x1234 may be marked with 01 (fixed).

(160) In some embodiments, the hardware logic block **900A** may be configured to use the metadata inputs to construct an input for a first rule cache lookup, such as a lookup for an access control rule. For instance, the hardware logic block **900A** may mask flow metadata from the source and/or the target of the data transfer instruction. Referring again to the example of FIG. 7B, the hardware logic block **900A** may mask the metadata symbol PTR BLUE associated with the data register **R1** and the metadata symbol PTR GREEN associated with the application memory address 0x1234, because both may be marked with 10 (flow). The resulting input to the rule cache **144** may include the following, where the symbol DEF denotes a default metadata value (e.g., a bit string of all 1's).

(161) **R0**=PTR RED, **R1**=DEF, mem={DEF, CEL RED}

(162) In some embodiments, the metadata {<<PTR GREEN>>*10, <<CEL RED>>*01} associated with the application memory address 0x1234 may be a 32-bit string, where the first 16 bits may encode PTR GREEN (with designated bits 10) and the last 16 bits may encode CEL RED (with designated bits 01). Thus, each of <<PTR GREEN>> and <<CEL RED>> may have 14 bits. To mask the metadata symbol PTR GREEN, the 32-bit string encoding {<<PTR GREEN>>*10, <<CEL RED>>*01} may be combined with a 32-bit mask via an OR operator, where the first 16 bits of the mask are all 1's, and the last 16 bits of the mask are all 0's. Although not shown in FIG. 9, the mask may, in some embodiments, be output by the tag map table **142**.

(163) It should be appreciated that aspects of the present disclosure are not limited any particular technique for masking flow metadata, or to masking flow metadata at all. In some embodiments, a default metadata value of all 0's may be used, and the 32-bit string encoding {<<PTR GREEN>>*10, <<CEL RED>>*01} may be combined with a 32-bit mask via an AND operator, where the first 16 bits of the mask are all 0's, and the last 16 bits of the mask are all 1's.

(164) In some embodiments, the hardware logic block **900B** may be configured to use the metadata inputs to construct an input for a second rule cache lookup, such as a lookup for an information flow rule. For instance, the hardware logic block **900B** may mask the metadata symbol PTR RED

associated with the address register **R0**, because metadata associated with an address register may not be relevant for an information flow rule. Additionally, or alternatively, the hardware logic block **900B** may mask fixed metadata from the source of the data transfer instruction. Additionally, or alternatively, the hardware logic block **900B** may mask both fixed metadata and flow metadata from the target of the data transfer instruction. The resulting input to the rule cache **144** may include the following, where the symbol DEF denotes a default metadata value (e.g., a bit string of all 1's).

(165) **R0=DEF, R1=PTR BLUE, mem=DEF**

(166) In this manner, the access control rule and the information flow rule may be enforced using separate rule cache lookups, which, as explained above, may reduce rule cache pressure. However, it should be appreciated that aspects of the present disclosure are not limited to enforcing different rules using separate rule cache lookups. In some embodiments, one or both of the access control rule and the information flow rule may be enforced without using a rule cache lookup (e.g., using one or more of the symbolic processing techniques described above in connection with FIGS. **6**, **8**). For instance, the rule cache lookup for the information flow rule may be replaced by hardware logic processing, as discussed above in connection with the illustrative sub-block **815** shown in FIG. **8**. Additionally, or alternatively, the hardware logic block **900A** may be configured to provide input to the illustrative fixed rule cache **144A**, as discussed above in connection with the illustrative sub-block **810** shown in FIG. **8**. Additionally, or alternatively, the hardware logic block **900A** may be configured to output the metadata label **PTR RED** associated with the address register **R0** and the metadata label **CEL RED** associated with the application memory address **0x1234** to the illustrative sub-block **610** shown in FIG. **6**, which may be adapted to check that the metadata label **PTR RED** and the metadata label **CEL RED** are content and location counterparts of the same color (which may be a form of partial equality), and if so, output the metadata label **PTR BLUE** provided by the hardware logic block **900B**.

(167) In some embodiments, a tag map table (e.g., the illustrative tag map table **142** shown in FIG. **1**) may be configured to allow multiple fetches of metadata on a single input application memory address. The fetched metadata may then be used to perform multiple lookups in a rule cache, for example, as described above in connection with FIG. **9**. Additionally, or alternatively, multiple rule caches may be used, for example, as described above in connection with FIG. **8**.

(168) FIG. **10** shows illustrative tag map table entries **1000A-B**, in accordance with some embodiments. For instance, the entries **1000A-B** may be in the illustrative tag map table **142** shown in FIG. **1**, and may be linked so that both entries may be accessed on a single input application memory address.

(169) In some embodiments, the tag map table **142** may be implemented using a fully associative array. Additionally, or alternatively, the tag map table **142** may be implemented using a hardware page table walker. If the fully associative array becomes full, one or more overflow entries may be stored in a table in a memory (e.g., the illustrative metadata memory **125** shown in FIG. **1**). If an input address does not match any entry in the fully associative array, the hardware page table walker may be used to walk through the table in the memory to identify an entry, if any, that matches the input address.

(170) In the example of FIG. **10**, the tag map table entry **1000A** stores, in an “Offset” field, a first offset value to be added to an input address in an application memory (e.g., the illustrative application memory **120** shown in FIG. **1**) to obtain a first address in a metadata memory (e.g., the illustrative metadata memory **125** shown in FIG. **1**). Similarly, the tag map table entry **1000B** may store, in an “Offset” field, a second offset value to be added to the input address to obtain a second metadata memory address. The first and second offset values are shown in FIG. **10**, respectively, as “Offset-A” and “Offset-B.”

(171) Additionally, or alternatively, the tag map table entry **1000A** may store information that associates the tag map table entry **1000A** with one or more application memory addresses. For

instance, the tag map table entry **1000A** may store a start address and an end address, which may indicate an address range with which the tag map table entry **1000A** is associated. In some embodiments, the same start address and end address may be stored in the tag map table entry **1000B**, although that is not required. In some embodiments, a different start address and/or a different end address may be stored in the tag map table entry **1000B** (e.g., a subrange of the address range of the tag map table entry **1000A**).

(172) In some embodiments, when a host processor (e.g., the illustrative host processor **110** shown in FIG. 1) attempts to load data from, or store data to, an application memory, a target address of the load/store instruction may be used by tag processing hardware (e.g., the illustrative tag processing hardware **140** shown in FIG. 1) to determine whether the load/store instruction should be allowed.

(173) In some embodiments, the tag processing hardware **140** may match the target address to an entry in the tag map table **142**, such as the tag map table entry **1000A**. For instance, the tag processing hardware **140** may determine that the target address falls within an application memory range indicated by the start address and the end address stored in the tag map table entry **1000A**. The tag processing hardware **140** may then add the offset value stored in the tag map table entry **1000A** (Offset-A) to the target address to obtain a first metadata memory address, and may use the first metadata memory address to retrieve, from the metadata memory **125**, a first metadata label associated with the target address.

(174) However, it should be appreciated that aspects of the present disclosure are not limited to any particular technique for mapping a target address to a metadata memory address. In some embodiments, a metadata memory address may be obtained by applying a mask to the target address. For instance, the mask may include N zeros, so that a block of $2^{\text{sup}.N}$ addresses may be mapped to a same metadata memory address. A result of masking the N least significant bits to zero may be compared to a first value. If the result matches the first value, the target address may be combined with a second value (e.g., using an OR operation) to obtain the metadata memory address. In some embodiments, the first value may be indicative of an address block of size $2^{\text{sup}.N}$ in the application memory, and the second value may be indicative of an offset with N trailing zeros. The first and second values may not share any power of 2, so that combining the target address with the second value using an OR operation may result in a sum of the target address and the offset. This may provide savings in area on a chip, because an OR operations may be simpler to implement than an addition operation.

(175) In some embodiments, the tag map table entries **1000A-B** may be linked together by setting a bit in a “Continue” field of the tag map table entry **1000A**. For instance, the tag processing hardware **140** may check the bit in the “Continue” field of the tag map table entry **1000A**. If that bit has been set to 1, the tag processing hardware **140** may continue to the tag map table entry **1000B**. The tag processing hardware **140** may add the offset value stored in the tag map table entry **1000B** (Offset-B) to the target address to obtain a second metadata memory address, and may use the second metadata memory address to retrieve, from the metadata memory **125**, a second metadata label associated with the target address.

(176) In some embodiments, the tag map table entry **1000B** may also have a “Continue” field that stores a bit. If that bit has been set to 1, the tag processing hardware **140** may continue to a next entry (not shown in FIG. 10) in the tag map table entry **142**. In this manner, any suitable number of entries in the tag map table **142** may be linked together, which may in turn allow any suitable number of policy rules to be enforced via separate rule cache lookups.

(177) However, it should be appreciated that aspects of the present disclosure are not limited to using a “Continue” field to link together multiple tag map table entries. In some embodiments, the tag processing hardware **140** may be configured to invoke the tag map table **142** multiple times using the same input application memory address. For instance, a counter may be used to keep track of a number of times the tag map table **142** is to be accessed.

(178) Additionally, or alternatively, the tag map table entry **1000A** may include a field (not shown in FIG. **10**) that stores a pointer pointing to another tag map table entry to be accessed, or to an additional tag map table storing one or more entries to be accessed. Additionally, or alternatively, the tag processing hardware **140** may be configured to decompose a binary representation of metadata (e.g., a 32-bit value) retrieved via a single tag map table fetch into multiple binary representations of metadata (e.g., first 16 bits as a first binary representation of metadata, and last 16 bits as a second binary representation of metadata) to be used in different rule cache lookups. (179) Moreover, while Offset-A and Offset-B in the example of FIG. **10** point to disparate regions of the metadata memory **125**, aspects of the present disclosure are not so limited. In some embodiments, Offset-A and Offset-B may point to a same region of the metadata memory **125**. For instance, Offset-A may point to an even byte, whereas Offset-B may point to an immediately following odd byte. In this manner, consecutive bytes may be retrieved, which may improve cache performance.

(180) Referring again to the examples shown in FIGS. **7B**, **9**, the target address matching the tag map table entry **1000A** may be the application memory address 0x1234 stored in the illustrative address register **R0**, whereas the first metadata retrieved from the first metadata memory address may include a binary representation of the metadata symbol CEL RED, and the second metadata retrieved from the second metadata memory address may include a binary representation of the metadata symbol PTR GREEN. In some embodiments, the first metadata may include one or more designated bits indicative of how the metadata symbol CEL RED is classified, such as 01 for fixed. Likewise, the second metadata may include one or more designated bits indicative of how the metadata symbol PTR GREEN is classified, such as 10 for flow.

(181) In some embodiments, the metadata memory **125** may include two regions, shown as **MMR1** and **MMR2** in the example of FIG. **10**. The region **MMR1** may be designated for storing fixed metadata, where the region **MMR2** may be designated for storing flow metadata. In this manner, the metadata symbols CEL RED and PTR GREEN may be stored without classification bits. For instance, one or more hardware logic blocks may be provided to strip the bits 01 before storing fixed metadata to the region **MMR1**, and/or add the bits 01 after retrieving fixed metadata from the region **MMR1**. Likewise, one or more hardware logic blocks may be provided to strip the bits 10 before storing flow metadata to the region **MMR2**, and/or add the bits 10 after retrieving fixed metadata from the region **MMR2**. Such hardware logic blocks may be part of, or separate from, the tag map table **142**.

(182) The inventors have recognized and appreciated that stripping classification bits prior to storage may reduce an amount of memory used for storing metadata. However, it should be appreciated that aspects of the present disclosure are not limited to storing metadata without classification bits.

(183) In the example shown in FIG. **10**, the application memory **120** and the metadata memory **125** are portions of a same physical memory. A memory management component may be used that implements an appropriate protection scheme to prevent instructions executing on the host processor from modifying metadata stored in the metadata memory **125**. However, it should be appreciated that aspects of the present disclosure are not so limited. In some embodiments, the application memory **120** and the metadata memory **125** may be implemented on physically separate memories, and an appropriate mapping scheme may be used to map application memory addresses to metadata memory addresses.

(184) In some embodiments, the tag map table entry **1000A** may include a field denoted “No. of Bits.” The No. of Bits field may store a value indicating a length of an encoded metadata value to be retrieved from the metadata memory. For instance, the No. of Bits field may store a value indicating that the encoded metadata value may have a length of zero bits, one bit, two bits, three bits, four bits, five bits, . . . , 8 bits (or one byte), . . . , 16 bits (or two bytes), etc. However, it should be appreciated that aspects of the present disclosure are not limited to indicating encoded

metadata length in terms of bits. In some embodiments, a “No. of Bytes” field may be provided. Similarly, the tag map table entry **1000B** may also include a field denoted “No. of Bits,” which may store a value indicating a length of an encoded metadata value to be retrieved from the metadata memory.

(185) The inventors have recognized and appreciated that the “No. of Bits” fields in the tag map table entries **1000A-B** may be configured with different values. For instance, in some embodiments, more distinct values may be used for fixed metadata compared to flow metadata (or vice versa). Accordingly, a higher value may be stored in the “No. of Bits” field in the tag map table entries **1000A** compared to the tag map table entries **1000B** (or vice versa). This may provide flexibility in metadata usage. However, it should be appreciated that aspects of the present disclosure are not limited to storing any particular value or combination of values in the “No. of Bits” fields.

(186) In some embodiments, the tag map table **142** may assign one or more classification bits to a region of the application memory **120**. For instance, the tag map table **142** may map an address in the application memory **120** to an address in the metadata memory **125**, and the one or more classification bits may be added to a tag retrieved from the metadata memory address.

(187) In some embodiments, the tag map table **142** may allow one or more classification bits to be stored in a smaller tag by changing a position of the one or more classification bits as the tag is loaded from the metadata memory **125**. For instance, an 8-bit tag in memory may be converted to a 32-bit system representation by shifting three classification bits from positions 7:5 in the 8-bit tag to positions 31:29 in the 32-bit system tag.

(188) Illustrative techniques for encoding/decoding metadata values are described in International Patent Application No.: PCT/US2019/060698, entitled “SYSTEMS AND METHODS FOR METADATA ENCODING,” filed on Nov. 11, 2019, which is incorporated herein by reference in its entirety. However, it should be appreciated that aspects of the present disclosure are not limited to any particular encoding/decoding technique.

(189) In some embodiments, the tag map table entry **1000A** may include a field denoted “Mask.” The Mask field may store a bit string to be combined with an encoding of metadata to mask one or more selected bits in the encoding. For instance, as described above in connection with FIG. 9, the tag map table **142** may output a 32-bit string encoding, {<<PTR GREEN>>*10, <<CEL RED>>*01}, of the metadata associated with the application memory address 0x1234, where the first 16 bits may encode PTR GREEN (with designated bits 10) and the last 16 bits may encode CEL RED (with designated bits 01). The Mask field of the tag map table entry **1000A** may store a 32-bit mask, where the first 16 bits of the mask are all 1's, and the last 16 bits of the mask are all 0's. The illustrative policy violation lookup **900A** shown in FIG. 9 may combine this mask with the 32-bit string encoding {<<PTR GREEN>>*10, <<CEL RED>>*01} to mask the metadata label PTR GREEN, and a result may be used to provide input to the rule cache **144**.

(190) Similarly, the tag map table entry **1000B** may include a field denoted “Mask” for storing one or more masking bits. For instance, the Mask field of the tag map table entry **1000B** may store a 32-bit mask, where the first 16 bits of the mask are all 0's, and the last 16 bits of the mask are all 1's. This mask may be combined with the 32-bit string encoding {<<PTR GREEN>>*10, <<CEL RED>>*01} to mask the metadata label CEL RED.

(191) However, it should be appreciated that aspects of the present disclosure are not limited to any particular technique for masking metadata, or to using a mask to construct a rule cache input or a metadata output.

(192) It should also be appreciated that aspects of the present disclosure are not limited to enforcing policies against instructions executed by a host processor. For instance, the inventors have recognized and appreciated that a transfer of data via direct memory access (DMA) may be viewed as a load instruction from a source memory location immediately followed by a store instruction to a target memory location, without storing the data in any intermediate data register. Accordingly, in

some embodiments, one or more of the techniques described herein may be used to check data transfer instructions performed by a DMA device.

(193) FIG. **11** shows an illustrative register file **1100**, in accordance with some embodiments. The register file **1100** may be a register file of a DMA device, and may include a plurality of registers, such as registers **R0**, **R1**, etc.

(194) In some embodiments, the DMA device may execute a data transfer instruction to copy data from a source location in an application memory (e.g., the illustrative application memory **120** shown in FIG. **1**) to a target location in the same application memory or a different application memory. An address (e.g., 0x1234) of the source location may be stored in a first address register (e.g., **R0**), and an address (e.g., 0x5678) of the target location may be stored in a second data register (e.g., **R1**).

(195) In the example of FIG. **11**, each register in the register file **1100** has a corresponding tag register for storing metadata. The corresponding tag registers may be in an illustrative tag register file **1146**. For instance, the register **R0** may have a corresponding tag register in the tag register file **1146**, and the corresponding tag register may store a binary representation <<PTR RED>> of a metadata label RED. Likewise, the register **R1** may have a corresponding tag register in the tag register file **1146**, and the corresponding tag register may store a binary representation <<PTR BLUE>> of a metadata label PTR BLUE.

(196) Additionally, or alternatively, the application memory address 0x1234 stored in the address register **R0** may have a corresponding address in a metadata memory (e.g., the illustrative metadata memory **125** shown in FIG. **1**). Metadata associated with the application memory address 0x1234 may be stored at the corresponding metadata memory address. For instance, in the example of FIG. **11**, a binary representation of metadata, {<<PTR GREEN>>, <<CEL RED>>}, may be stored at the metadata memory address corresponding to the application memory address 0x1234. Similarly, the application memory address 0x5678 stored in the address register **R1** may have a corresponding address in the metadata memory **125**. Metadata associated with the application memory address 0x5678 may be stored at the corresponding metadata memory address. For instance, in the example of FIG. **11**, a binary representation of metadata, {<<PTR YELLOW>>, <<CEL BLUE>>}, may be stored at the metadata memory address corresponding to the application memory address 0x5678.

(197) In some embodiments, tag processing hardware (e.g., the illustrative tag processing hardware **140** shown in FIG. **1**, or a separate instance of tag processing hardware for checking DMA instructions) may apply one or more rules to determine whether a DMA instruction should be allowed to copy data stored at the application memory address 0x1234 to the application memory address 0x5678. Such a rule may be provided by a policy processor and/or a rule cache (e.g., the illustrative policy processor **150** and/or the illustrative rule cache **144** shown in FIG. **1**).

Additionally, or alternatively, symbolic rule processing may be used, for example, as described above in connection with FIGS. **5-6**.

(198) For example, an access control rule may be provided that checks whether flow metadata associated with the source address register **R0** matches fixed metadata associated with the source location referenced by the application memory address 0x1234 stored in the source address register **R0**. In the example of FIG. **11**, the metadata label PTR RED associated with the source address register **R0** may be marked with 10 (flow), whereas the metadata label CEL RED associated with the application memory address 0x1234 may be marked with 01 (fixed).

(199) Additionally, or alternatively, the access control rule may check whether flow metadata associated with the target address register **R1** matches fixed metadata associated with the target location referenced by the application memory address 0x5678 stored in the target address register **R0**. In the example of FIG. **11**, the metadata label PTR BLUE associated with the target address register **R1** may be marked with 10 (flow), whereas the metadata label CEL BLUE associated with the application memory address 0x5678 may be marked with 01 (fixed).

(200) One or more of the techniques described herein may be used to apply an access control rule

against metadata labels. For instance, a rule cache lookup may be performed (e.g., using one or more of the techniques described above in connection with FIGS. **8-10**), and/or a hardware logic block may be provided to check whether two input metadata labels (e.g., PTR RED and CEL RED) are counterparts of the same color (e.g., using one or more of the techniques described above in connection with FIGS. **5-6**).

(201) In some embodiments, an information flow rule for DMA may be provided. The inventors have recognized and appreciated that, in a DMA context, there may be no metadata update for the source location. By contrast, for the target location, fixed metadata may be preserved, whereas flow metadata may be updated (e.g., replaced by flow metadata from the source location). One or more of the techniques described herein may be used to apply such an information flow rule. For instance, a rule cache lookup may be performed (e.g., using one or more of the techniques described above in connection with FIGS. **8-10**), and/or a hardware logic block may be provided to forward flow metadata without decoding (e.g., using one or more of the techniques described above in connection with FIGS. **5-6**).

(202) The inventors have further recognized and appreciated that, in some embodiments, one or more classification bits may be used to allow a rule cache to be omitted in favor of a small amount of hardware logic (e.g., the illustrative hardware accelerator **500** shown in FIG. **5**). If the hardware logic encounters a combination of one or more input signals that the hardware logic is unable to process, the hardware logic may invoke a policy processor (e.g., the illustrative policy processor **150** shown in FIG. **1**). In this manner, chip area may be reduced (e.g., by omitting the rule cache), and/or metadata processing speed may be improved on common inputs.

(203) It should be appreciated that DMA transfers are not limited to memory-to-memory transfers. The techniques described herein for protecting DMA transfers may, additionally or alternatively, be used to protect other types of DMA transfers (e.g., to and/or from a peripheral device such as a USB storage device).

(204) The inventors have also recognized and appreciated that, to check an instruction of a certain type, more than three input metadata labels may be presented to a rule cache. By contrast, many common instructions have only three input metadata labels (e.g., **R0**, **R1**, and **mem** in the examples discussed above in connection with FIGS. **3-10**). Thus, memory space may be wasted by a rule cache that is able to accommodate four inputs.

(205) For instance, some instruction sets have vector instructions, and a vector instruction may perform a plurality of operations across multiple memory locations over multiple clock cycles. In some embodiments, each such memory location may be associated with a respective metadata label, so that there may be heterogeneous metadata across the vector. In some embodiments, one or more classification bits (e.g., 00 for neither, 01 for fixed, 10 for flow, 11 for both) may be provided for the respective metadata labels, and one or more of the techniques described herein may be used to check the vector instruction against one or more policies.

(206) In some embodiments, a vector instruction may include a multiply-accumulate (MAC) instruction, which may multiply each element of a vector with a scalar value, and accumulate the results. Three registers may be used: an address register holding an application memory address at which a current element of the vector is stored, a scalar register holding the scalar value, and an accumulator register for storing a result that has been accumulated thus far. A host processor (e.g., the illustrative host processor **110** in the example of FIG. **1**) may step through the vector in a number of iterations. At each iteration, the host processor **110** may use the application memory address stored in the address register to retrieve the current element of the vector, multiply the current element with the scalar value stored in the scalar register, add the multiplication result to the current accumulation result stored in the accumulator register to obtain a new accumulation result, and store the new accumulation result in the accumulator register.

(207) Thus, to check a MAC instruction, four input metadata labels may be presented to a rule cache, including a respective metadata label associated with each of the three registers, and a

metadata label associated with the location referenced by the application memory address stored in the address register.

(208) Accordingly, in some embodiments, one or more of the techniques described herein may be used to check a MAC instruction via multiple rule cache lookups. Each lookup may have fewer than four inputs, so that a narrower rule cache may be used. For instance, a first lookup may be performed to apply an access control rule, for example, to determine whether the MAC instruction should be allowed. The access control may be applied against flow metadata associated with the address register and fixed metadata associated with the location referenced by the application memory address stored in the address register. Additionally, or alternatively, a second lookup may be performed to apply an information flow rule, for example, to update flow metadata associated with the accumulator register based on flow metadata associated with the scalar register and/or flow metadata associated with the application memory address stored in the address register.

(209) In some embodiments, one or both of these rule cache lookups may be replaced by processing performed by one or more hardware logic blocks, for example, as discussed above in connection with the examples of FIGS. 5-6, 8-10. For instance, if no input metadata label is marked with 10 (flow), the second rule cache lookup may be omitted, thereby reducing rule cache bandwidth by 50%.

(210) It should be appreciated that vector instructions such as MAC are discussed above solely for purposes of illustration. The techniques described herein may be used to check any suitable instruction in addition to, or instead of, vector instructions. For instance, in some embodiments, PUSH and/or POP instructions which transfer multiple values between an in-memory stack and a register file may be checked using one or more of the techniques described herein.

(211) Thus, the inventors have recognized and appreciated that fixed and/or flow classification bits may be used to allow certain checks to be performed by one or more simple hardware circuits, so that a narrower rule cache may be used. Additionally, or alternatively, such classification bits may be used to improve a performance of a wider rule cache. However, it should be appreciated that aspects of the present disclosure are not limited to using any particular types of classification bits, or any classification bit at all.

(212) In some embodiments, one or more classification bits may be used to provide differentiated processing of metadata. For instance, one or more designated bits in a binary representation of metadata may indicate how a remaining portion of the binary representation should be parsed (e.g., by the illustrative hardware accelerator 500 in the example of FIG. 5). As an example, a 32-bit binary representation may include 4 designated bits which, when interpreted, may cause the hardware accelerator 500 to parse the remaining 28 bits into an 18-bit chunk and a 10-bit chunk. The 18-bit chunk may be interpreted as a 16-bit encoding of metadata augmented with two classification bits, whereas the 10-bit chunk may be interpreted as an 8-bit encoding of metadata augmented with two classification bits. The 16-bit encoding may represent a set of two metadata symbols (e.g., {A, B}), while the 8-bit encoding may represent a set of one metadata symbol (e.g., {C}).

(213) Thus, the one or more designated bits may indicate how a set of metadata symbols (e.g., {A, B, C}) may be divided into two or more of subsets (e.g., {A, B} and {C}). However, it should be appreciated that aspects of the present disclosure are not limited to parsing metadata in any particular manner, or at all.

(214) In some embodiments, a hardware accelerator may process a chunk of metadata depending on one or more associated classification bits. For instance, referring to the above example, the two classification bits for the 16-bit encoding may indicate that the 16-bit encoding represents two metadata symbols to be passed along to a rule cache and/or a policy processor (e.g., the illustrative rule cache 144 and/or the illustrative policy processor 150 in the example of FIG. 1). Additionally, or alternatively, the classification bits for the 8-bit encoding may indicate that the 8-bit encoding represents a metadata symbol to be processed symbolically (e.g., as described above in connection

with the example of FIG. 6).

(215) FIG. 12 shows an illustrative hardware logic block **1200**, in accordance with some embodiments. The hardware logic block **1200** may be similar to the illustrative hardware logic block **900A** in the example of FIG. 9. For instance, the hardware logic block **1200** may be configured to look up the illustrative rule cache **144** based on one or more metadata inputs. Additionally, or alternatively, the hardware logic block **1200** may provide a comparison output to the rule cache **144**.

(216) Similar to the hardware logic block **900A**, the hardware logic block **1200** may receive three metadata inputs from tag fetch logic, which may include the tag map table **142**. Referring to the example of FIG. 7B, the first input may be <<PTR RED>>*10, the metadata read from the tag register corresponding to the address register **R0**, whereas the second input may be <<PTR BLUE>>*10, the metadata read from the tag register corresponding to the data register **R1**, and the third input may be {<<PTR GREEN>>*10, <<CEL RED>>*01}, the metadata read from the metadata memory address corresponding to the application memory address 0x1234. The hardware logic block **1200** may use classification bits to identify flow metadata associated with the address register **R0** (e.g., <<PTR RED>>) and fixed metadata associated with the application memory address 0x1234 (e.g., <<CEL RED>>).

(217) In some embodiments, the binary representation <<PTR RED>> may include a first portion corresponding to the constructor PTR and a second portion corresponding to the color RED. This may be denoted <<PTR>><<RED>>. Likewise, the binary representation <<CEL RED>> may include a first portion corresponding to the constructor CEL and a second portion corresponding to the color RED. This may be denoted <<CEL>><<RED>>. The hardware logic block **1200** may be configured to parse such metadata to identify color portions. Additionally, or alternatively, the hardware logic block **1200** may be configured to perform a comparison on the color portions (e.g., using one or more of the techniques described above in connection with FIGS. 5-6).

(218) In some embodiments, the hardware logic block **1200** may provide a comparison result to the rule cache **144**, which may process one or more rule cache inputs based on the comparison result. For instance, the rule cache **144** may use classification bits to identify flow metadata associated with the address register **R0** (e.g., <<PTR RED>>) and fixed metadata associated with the application memory address 0x1234 (e.g., <<CEL RED>>). In response to identifying such metadata, the rule cache **144** may check the comparison result received from the hardware logic block **1200**.

(219) If the comparison result indicates there is a match in color, the rule cache **144** may, before performing a lookup, mask color portions of the flow metadata associated with the address register **R0** and/or the fixed metadata associated with the application memory address 0x1234 (e.g., using one or more of the techniques described above in connection with FIG. 9). For instance, the rule cache **144** may replace <<RED>> in <<PTR>><<RED>> and <<CEL>><<RED>> with some default metadata <<DEF>>, resulting in <<PTR>><<DEF>> and <<CEL>><<DEF>>, respectively.

(220) Additionally, or alternatively, the hardware logic block **1200** may use classification bits to mask color portions of the flow metadata associated with the address register **R0** and/or the fixed metadata associated with the application memory address 0x1234 (e.g., using one or more of the techniques described above in connection with FIG. 9). The rule cache **144** may perform a lookup based on masked inputs (e.g., <<PTR>><<DEF>> and <<CEL>><<DEF>>). Another hardware logic block may be provided to combine an output of the rule cache **144** and the comparison result received from the hardware logic block **1200**.

(221) In this manner, instead of storing an entry for each color value, the rule cache **144** may simply store a single entry, such as the following. 1) If **R0**=PTR DEF and mem=CEL DEF, then allow store instruction

(222) Thus, concrete metadata values may be abstracted away from rule cache entries, and may

instead be processed symbolically by the hardware logic block **1200**. This may reduce rule cache pressure by a factor of N (the number of distinct color values). However, it should be appreciated that aspects of the present disclosure are not limited to masking color portions of metadata presented to a rule cache, or to having separate color portions at all.

(223) The inventors have recognized and appreciated that, in some instances, determining whether an instruction is to be allowed may be done more expeditiously than determining updated metadata. For instance, determining whether to allow an instruction may simply involve comparing metadata symbols, which may be performed by a hardware logic block (e.g., as described above in connection with the example of FIG. 6). By contrast, determining updated metadata may involve looking up the rule cache **144** and/or invoking the policy processor **150** (e.g., for a DMA transfer).

(224) The inventors have recognized and appreciated that it may be desirable to allow a result of executing a first instruction to be written back to memory (e.g., the illustrative application memory **120** in the example of FIG. 1) as soon as it is determined that the first instruction is to be allowed, even if updated metadata is not yet available. For example, this may allow a host processor (e.g., the illustrative host processor **110** in the example of FIG. 1) to proceed with one or more second instructions, while the updated metadata for the first instruction is still being looked up and/or computed.

(225) However, the inventors have recognized and appreciated that, if a result of executing a first instruction is written to an application memory location before corresponding metadata is available, that result may be accessed by a second instruction before a metadata memory location corresponding to the application memory location is updated. As a result, the second instruction may be checked for policy violation based on out-of-date metadata stored at the corresponding metadata memory location, which may lead to failure to detect a policy violation.

(226) Accordingly, in some embodiments, one or more designated bits may be provided in a binary representation of metadata to indicate whether a portion of the metadata is out of date. For instance, a tag register associated with a register storing a result of executing an instruction may store a binary representation of metadata that includes a first portion and a second portion, along with a classification bit for the first portion to indicate whether the first portion is out of date, and likewise for the second portion. Such a classification bit may be referred to herein as an out-of-date bit.

(227) In some embodiments, the binary representation may further include two classification bits for the first portion indicating that the first portion includes fixed metadata, as well as two classification bits for the second portion indicating that the second portion includes flow metadata. Thus, when the result of executing the instruction is stored into the register, only the second portion of metadata (flow) may be updated.

(228) In some embodiments, if it is determined that the instruction is to be allowed, but updated metadata for the second portion is not yet available, the out-of-date bit for the second portion of metadata (flow) may be set to 1. When the updated metadata becomes available (e.g., as returned by the rule cache **144** and/or the policy processor **150**), the second portion of metadata may be updated, and the corresponding out-of-date bit may be set to 0.

(229) In this manner, when it is determined that the instruction is to be allowed, the result of executing the instruction may be written to an application memory location, even if updated metadata is not yet available. For instance, a store instruction attempting to write the result to the application memory location may be checked based on the first portion of metadata associated with the result register (fixed), without regard to the second portion of metadata associated with the result register (flow). Assuming the out-of-date bit for the first portion of metadata remains 0, the store instruction may be allowed. As a result, the second portion of metadata, along with the associated out-of-date bit (which has been set to 1), may be carried over to a metadata memory location corresponding to the application memory location.

(230) In some embodiments, tag processing hardware (e.g., by the illustrative tag processing hardware **140** in the example of FIG. 1) may be configured to check an out-of-date bit for a portion

of metadata before using that portion of metadata to check policy compliance. Thus, if a subsequent instruction attempts to access the result from the application memory location while the out-of-date bit is still set to 1, the tag processing hardware **140** may determine that the application memory location is associated with out-of-date metadata, and may wait for the updated metadata to become available before checking the subsequent instruction.

(231) In some embodiments, the tag processing hardware **140** may maintain a record of registers and/or application memory locations that are associated with out-of-date metadata. Such a record may include dependency information (e.g., one or more data transfer instructions involving out-of-date flow metadata). Thus, when the updated metadata becomes available for the result register, the tag processing hardware **140** may use the dependency information to map the result register to the application memory location at which the result has been stored, and may use the updated metadata to update the metadata memory location corresponding to the application memory location.

(232) Additionally, or alternatively, the tag processing hardware **140** may maintain a record of classification bits associated with one or more registers and/or application memory locations. Such a record may be updated based on data dependencies among the one or more registers and/or application memory locations. For instance, for an arithmetic instruction with first and second operands, an OR operation may be performed on first and second flow bits that are associated, respectively, with registers storing the first and second operands. A result of the OR operation may be used to update a flow bit associated with a register storing a result of the arithmetic instruction. In this manner, one or more logical operations may be used to maintain an up-to-date record of classification bits for metadata being processed by the tag processing hardware **140**.

(233) The inventors have recognized and appreciated various advantages of maintaining an up-to-date record of classification bits. For instance, as described above in connection with FIG. **8**, classification bits may be used to determine whether certain instructions may be quickly disposed of without further metadata processing (e.g., arithmetic instructions on application data with neither flow metadata nor fixed metadata). Therefore, a record of classification bits may be used to eliminate instructions to be checked by the tag processing hardware **140**.

(234) Illustrative configurations of various aspects of the present disclosure are provided below.

(235) A1. A method for metadata processing, the method comprising acts of: receiving one or more metadata inputs; and processing the one or more metadata inputs to determine whether to allow an instruction, comprising: identifying one or more classification bits from a metadata input of the one or more metadata inputs; and processing the metadata input based on the one or more classification bits.

(236) A2. The method of configuration A1, wherein: the one or more classification bits are identified without decoding the metadata input.

(237) A3. The method of configuration A1, wherein: processing the metadata input based on the one or more classification bits comprises: parsing, based on the one or more classification bits, the metadata input into at least a first portion and a second portion.

(238) A4. The method of configuration A3, wherein: processing the metadata input based on the one or more classification bits further comprises: processing the first portion and the second portion of the metadata input differently.

(239) A5. The method of configuration A4, wherein: the one or more classification bits comprise one or more first classification bits; and the method further comprises: identifying one or more second classification bits from the first portion of the metadata input; and processing the first portion of the metadata input based on the one or more second classification bits.

(240) A6. The method of configuration A1, wherein: processing the metadata input based on the one or more classification bits comprises: determining, based on the one or more classification bits, whether to perform a rule cache lookup based on the metadata input.

(241) A7. The method of configuration A1, wherein: processing the metadata input based on the one or more classification bits comprises: determining, based on the one or more classification bits,

whether to use at least one hardware logic block to process the metadata input.

(242) A8. The method of configuration A7, wherein: the metadata input comprises a first metadata input; the one or more metadata inputs further comprises a second metadata input; and the at least one hardware logic block is configured to perform a comparison between a portion of the first metadata input and a portion of the second metadata input.

(243) A9. The method of configuration A7, wherein: the at least one hardware logic block is programmable to perform different types of processing on the metadata input.

(244) A10. The method of configuration A7, wherein: processing the metadata input based on the one or more classification bits further comprises: selecting, based on the one or more classification bits, the at least one hardware logic block from a plurality of hardware logic blocks.

(245) A11. The method of configuration A10, wherein: the plurality of hardware logic blocks correspond, respectively, to a plurality of policies; and the method further comprises an act of: activating the at least one hardware logic block in response to at least one corresponding policy being loaded.

(246) A12. The method of configuration A7, wherein: the at least one hardware logic block is configured to construct, based on the metadata input, respective inputs for a plurality of rule cache lookups.

(247) A13. The method of configuration A12, wherein: the instruction comprises a vector instruction having a plurality of component operations; and the plurality of rule cache lookups correspond, respectively, to the plurality of component operations.

(248) A14. The method of configuration A12, wherein: the instruction comprises a direct memory access instruction having at least one load component and at least one store component; and the plurality of rule cache lookups comprise at least first and second lookups corresponding, respectively, to the at least one load component and the at least one store component.

(249) A15. The method of configuration A1, wherein: the metadata input is associated with a register or memory location used by the instruction; and the one or more classification bits indicate the metadata input is for determining whether the instruction is allowed to use the register or memory location.

(250) A16. The method of configuration A15, wherein: processing the metadata input based on the one or more classification bits comprises: using the metadata input to construct an input for looking up a rule cache storing one or more allowed metadata input patterns.

(251) A17. The method of configuration A1, wherein: the metadata input is associated with a first register or memory location used by the instruction; and the one or more classification bits indicate the metadata input is for performing a metadata update for a second register or memory location used by the instruction.

(252) A18. The method of configuration A17, wherein: processing the metadata input based on the one or more classification bits comprises: using the metadata input to construct an input for looking up a rule cache storing one or more entries that map metadata input patterns to metadata output patterns.

(253) B1. A method for metadata processing, the method comprising acts of: receiving an application memory address; accessing mapping information associated with the application memory address; and using the mapping information to retrieve metadata from a memory region of a plurality of metadata memory regions, wherein the plurality of metadata memory regions correspond, respectively, to a plurality of metadata classifications.

(254) B2. The method of configuration B1, wherein: the metadata retrieved from the memory region is in a representation for storage; and the method further comprises an act of:

(255) transforming the metadata into a representation for processing, wherein the representation for processing comprises one or more classification bits based on the metadata classification of the metadata memory region from which the metadata is retrieved.

(256) B3. The method of configuration B2, wherein: transforming the metadata comprises

augmenting the representation for storage with the one or more classification bits based on the metadata classification of the metadata memory region from which the metadata is retrieved.

(257) B4. The method of configuration B2, wherein: the representation for storage has a first length; the representation for processing has a second length different from the first length; transforming the metadata comprises shifting the one or more classification bits from a first position in the representation for storage to a second position in the representation for processing.

(258) B5. The method of configuration B2, wherein: the one or more classification bits indicate the metadata is for determining whether an instruction is allowed to use a memory location identified by the application memory address.

(259) B6. The method of configuration B2, wherein: the application memory address identifies a first memory location used by an instruction; and the one or more classification bits indicate the metadata is for performing a metadata update for a second memory location or a register used by the instruction.

(260) B7. The method of configuration B1, wherein: the mapping information comprises first mapping information; the metadata comprises first metadata; the memory region comprises a first memory region; and the method further comprises acts of: accessing second mapping information associated with the application memory address; and using the second mapping information to retrieve second metadata from a second memory region of the plurality of metadata memory regions.

(261) B8. The method of configuration B7, wherein: the metadata classification of the first metadata memory region comprises a first metadata classification; the first metadata classification indicates the first metadata is for determining whether an instruction is allowed to use a memory location identified by the application memory address; and the method further comprises an act of: using the first metadata to construct an input for looking up a rule cache storing one or more allowed metadata input patterns.

(262) B9. The method of configuration B7, wherein: the application memory address identifies a first memory location used by an instruction; the metadata classification of the second metadata memory region comprises a second metadata classification; the second metadata classification indicates the second metadata is for performing a metadata update for a second register or memory location used by the instruction; and the method further comprises an act of: using the second metadata to construct an input for looking up a rule cache storing one or more entries that map metadata input patterns to metadata output patterns.

(263) B10. The method of configuration B7, wherein: the first mapping information is accessed from a first entry in a tag map table; the first entry includes a continue bit; and the method further comprises an act of: determining, based on the continue bit in the first entry, whether to access another entry in the tag map table.

(264) B11. The method of configuration B10, wherein: the second mapping information is accessed from a second entry in the tag map table in response to determining that another entry is to be accessed.

(265) C1. A method for metadata processing, the method comprising acts of: identifying at least one metadata symbol from a metadata label; identifying at least one attribute value associated with the at least one metadata symbol; using the at least one attribute value to determine one or more classification bits; generating a binary representation for the metadata label, wherein: the binary representation comprises at least one portion representing the at least one metadata symbol; and the binary representation associates the one or more classification bits with the at least one portion representing the at least one metadata symbol.

(266) C2. The method of configuration C1, wherein: the at least one attribute value is identified based on at least one data type of the at least one metadata symbol.

(267) C3. The method of configuration C1, wherein: the at least one metadata symbol comprises a constructor and one or more arguments; and identifying at least one attribute value associated with

the at least one metadata symbol comprises: identifying a first attribute value associated with the constructor; and/or identifying a second attribute value associated with an argument of the one or more arguments.

(268) C4. The method of configuration C3, wherein: the first attribute value indicates the at least one metadata symbol is for determining whether an instruction is allowed.

(269) C5. The method of configuration C3, wherein: the first attribute value indicates the at least one metadata symbol is for performing metadata update.

(270) C6. The method of configuration C3, wherein: the second attribute value indicates the argument is for comparison against another value of a same data type as the argument.

(271) C7. The method of configuration C6, wherein: the data type of the argument comprises a set of color values; and the method further comprises acts of: conducting test runs of at least one software application, the at least one software application comprising a plurality of software components; assigning, to each of the plurality of software components, a color value from the set of color values; collecting statistics on the color values from the test runs of the at least one software application; and reassigning, based on the collected statistics, the color values to the plurality of software components.

(272) C8. The method of configuration C3, wherein: the at least one attribute value comprises the first attribute value and the second attribute value; and using the at least one attribute value to determine one or more classification bits comprises encoding the first and second attribute values into the one or more classification bits.

(273) D1. A method for metadata processing, the method comprising acts of: identifying one or more registers and/or one or more application memory locations used by an instruction; accessing one or more classification bits for the one or more registers and/or one or more application memory locations; and determining, based on the one or more classification bits, whether to allow the instruction.

(274) D2. The method of configuration D1, wherein: the one or more classification bits are accessed from a record of classification bits for a plurality of registers and/or application memory locations; the one or more registers and/or one or more application memory locations used by the instruction comprise one or more input registers and/or one or more input application memory locations; and the method further comprises an act of: in response to determining that the instruction is allowed, updating at least one entry in the record, the at least one entry corresponding to one or more output registers and/or one or more output application memory locations used by the instruction.

(275) D3. The method of configuration D2, wherein: the one or more input registers and/or one or more input application memory locations comprise a first input register or application memory location and a second input register or application memory location; the one or more classification bits comprise one or more first classification bits for the first input register or application memory location; the one or more classification bits further comprise one or more second classification bits for the second input register or application memory location; and updating the at least one entry comprises acts of: performing at least one logical operation on the one or more first classification bits and/or the one or more second classification bits to obtain one or more third classification bits; and storing the one or more third classification bits in the at least one entry.

(276) D4. The method of configuration D1, wherein: the instruction is determined to be allowed based on the one or more classification bits, without performing a rule cache lookup or invoking a policy processor.

(277) D5. The method of configuration D1, wherein: the instruction is not determined to be allowed based on the one or more classification bits, and the method further comprises an act of: accessing additional metadata associated with the one or more registers and/or one or more application memory locations; and performing a rule cache lookup and/or invoking a policy processor based on the additional metadata to determine if the instruction is allowed.

(278) E1. A system comprising circuitry and/or one or more processors programmed by executable instructions, wherein the circuitry and/or the one or more programmed processors are configured to perform the method of any of configurations A1-A18, B1-B11, C1-C8.

(279) E2. At least one computer-readable medium having stored thereon at least one netlist for the circuitry of configuration E1.

(280) E3. At least one computer-readable medium having stored thereon at least one hardware description that, when synthesized, produces the at least one netlist of configuration E2.

(281) E4. The at least one computer-readable medium of configuration E3, wherein the at least one hardware description is in an encrypted form.

(282) E5. At least one computer-readable medium having stored thereon the executable instructions of configuration E1.

(283) FIG. **13** shows, schematically, an illustrative computer **1000** on which any aspect of the present disclosure may be implemented.

(284) In the example shown in FIG. **13**, the computer **1000** includes a processing unit **1001** having one or more processors and a non-transitory computer-readable storage medium **1002** that may include, for example, volatile and/or non-volatile memory. The memory **1002** may store one or more instructions to program the processing unit **1001** to perform any of the functions described herein. The computer **1000** may also include other types of non-transitory computer-readable medium, such as storage **1005** (e.g., one or more disk drives) in addition to the system memory **1002**. The storage **1005** may also store one or more application programs and/or resources used by application programs (e.g., software libraries), which may be loaded into the memory **1002**.

(285) The computer **1000** may have one or more input devices and/or output devices, such as devices **1006** and **1007** illustrated in FIG. **13**. These devices may be used, for instance, to present a user interface. Examples of output devices that may be used to provide a user interface include printers and display screens for visual presentation of output, and speakers and other sound generating devices for audible presentation of output. Examples of input devices that may be used for a user interface include keyboards and pointing devices (e.g., mice, touch pads, and digitizing tablets). As another example, the input devices **1007** may include a microphone for capturing audio signals, and the output devices **1006** may include a display screen for visually rendering, and/or a speaker for audibly rendering, recognized text.

(286) In the example shown in FIG. **13**, the computer **1000** also includes one or more network interfaces (e.g., the network interface **1010**) to enable communication via various networks (e.g., the network **1020**). Examples of networks include a local area network (e.g., an enterprise network) and a wide area network (e.g., the Internet). Such networks may be based on any suitable technology and operate according to any suitable protocol, and may include wireless networks and/or wired networks (e.g., fiber optic networks).

(287) Having thus described several aspects of at least one embodiment, it is to be appreciated that various alterations, modifications, and improvements will readily occur to those skilled in the art. Such alterations, modifications, and improvements are intended to be within the spirit and scope of the present disclosure. Accordingly, the foregoing descriptions and drawings are by way of example only.

(288) The above-described embodiments of the present disclosure may be implemented in any of numerous ways. For example, the embodiments may be implemented using hardware, software, or a combination thereof. When implemented in software, the software code may be executed on any suitable processor or collection of processors, whether provided in a single computer, or distributed among multiple computers.

(289) Also, the various methods or processes outlined herein may be coded as software that is executable on one or more processors running any one of a variety of operating systems or platforms. Such software may be written using any of a number of suitable programming languages and/or programming tools, including scripting languages and/or scripting tools. In some instances,

such software may be compiled as executable machine language code or intermediate code that is executed on a framework or virtual machine. Additionally, or alternatively, such software may be interpreted.

(290) The techniques disclosed herein may be embodied as a non-transitory computer-readable medium (or multiple computer-readable media) (e.g., a computer memory, one or more floppy discs, compact discs, optical discs, magnetic tapes, flash memories, circuit configurations in Field Programmable Gate Arrays or other semiconductor devices, or other non-transitory, tangible computer storage medium) encoded with one or more programs that, when executed on one or more processors, perform methods that implement the various embodiments of the present disclosure discussed above. The computer-readable medium or media may be transportable, such that the program or programs stored thereon may be loaded onto one or more different computers or other processors to implement various aspects of the present disclosure as discussed above.

(291) The terms “program” or “software” are used herein to refer to any type of computer code or set of computer-executable instructions that may be employed to program one or more processors to implement various aspects of the present disclosure as discussed above. Moreover, it should be appreciated that according to one aspect of this embodiment, one or more computer programs that, when executed, perform methods of the present disclosure need not reside on a single computer or processor, but may be distributed in a modular fashion amongst a number of different computers or processors to implement various aspects of the present disclosure.

(292) Computer-executable instructions may be in many forms, such as program modules, executed by one or more computers or other devices. Program modules may include routines, programs, objects, components, data structures, etc. that perform particular tasks or implement particular abstract data types. Functionalities of the program modules may be combined or distributed as desired in various embodiments.

(293) Also, data structures may be stored in computer-readable media in any suitable form. For simplicity of illustration, data structures may be shown to have fields that are related through location in the data structure. Such relationships may likewise be achieved by assigning storage for the fields to locations in a computer-readable medium that convey relationship between the fields. However, any suitable mechanism may be used to establish a relationship between information in fields of a data structure, including through the use of pointers, tags, or other mechanisms that establish relationship between data elements.

(294) Various features and aspects of the present disclosure may be used alone, in any combination of two or more, or in a variety of arrangements not specifically discussed in the embodiments described in the foregoing, and are therefore not limited to the details and arrangement of components set forth in the foregoing description or illustrated in the drawings. For example, aspects described in one embodiment may be combined in any manner with aspects described in other embodiments.

(295) Also, the techniques disclosed herein may be embodied as methods, of which examples have been provided. The acts performed as part of a method may be ordered in any suitable way. Accordingly, embodiments may be constructed in which acts are performed in an order different from illustrated, which may include performing some acts simultaneously, even though shown as sequential acts in illustrative embodiments.

(296) Use of ordinal terms such as “first,” “second,” “third,” etc., in the claims to modify a claim element does not by itself connote any priority, precedence, or order of one claim element over another or the temporal order in which acts of a method are performed, but are used merely as labels to distinguish one claim element having a certain name from another element having a same name (but for use of the ordinal term) to distinguish the claim elements.

(297) Also, the phraseology and terminology used herein is for the purpose of description and should not be regarded as limiting. The use of “including,” “comprising,” “having,” “containing,”

“involving,” and variations thereof herein, is meant to encompass the items listed thereafter and equivalents thereof as well as additional items.

Claims

1. A method for metadata processing, the method comprising acts of: receiving one or more metadata inputs associated with an instruction, wherein: the one or more metadata inputs comprises a metadata input from a metadata storage location corresponding to an application data storage location used by the instruction; the metadata input comprises one or more first bit positions and one or more second bit positions; the one or more first bit positions of the metadata input encode at least one metadata symbol; and the one or more second bit positions of the metadata input comprise one or more classification bits indicating one or more attribute values associated with the at least one metadata symbol; and processing the one or more metadata inputs to determine whether to allow the instruction, comprising: identifying the one or more classification bits from the metadata input; and processing the metadata input based on the one or more classification bits.
2. The method of claim 1, wherein: the one or more classification bits are identified without decoding the one or more first bit positions of the metadata input that encode the at least one metadata symbol.
3. The method of claim 1, wherein: the at least one metadata symbol comprises at least one first metadata symbol and at least one second metadata symbol; and processing the metadata input based on the one or more classification bits comprises: parsing, based on the one or more classification bits, the one or more first bit positions of the metadata input into at least a first portion and a second portion encoding, respectively, the at least one first metadata symbol and the at least one second metadata symbol.
4. The method of claim 3, wherein: processing the metadata input based on the one or more classification bits further comprises: processing the first portion of the metadata input in a first manner; and processing the second portion of the metadata input in a second manner different from the first manner.
5. The method of claim 4, wherein: the one or more classification bits comprise one or more first classification bits; and the method further comprises: identifying one or more second classification bits from the first portion of the metadata input; and processing the first portion of the metadata input based on the one or more second classification bits.
6. The method of claim 1, wherein: processing the metadata input based on the one or more classification bits comprises: determining, based on the one or more classification bits, whether to perform a rule query based on the metadata input.
7. The method of claim 1, wherein: processing the metadata input based on the one or more classification bits comprises: determining, based on the one or more classification bits, whether to use at least one hardware logic block to process the metadata input.
8. The method of claim 7, wherein: the metadata input comprises a first metadata input; the one or more metadata inputs further comprises a second metadata input; and the at least one hardware logic block is configured to perform a comparison between a portion of the first metadata input and a portion of the second metadata input.
9. The method of claim 7, wherein: the at least one hardware logic block is programmable at run time to perform different types of processing on the metadata input.
10. The method of claim 7, wherein: determining, based on the one or more classification bits, whether to use at least one hardware logic block comprises: selecting, based on the one or more classification bits, the at least one hardware logic block from a plurality of hardware logic blocks.
11. The method of claim 10, wherein: the plurality of hardware logic blocks correspond, respectively, to a plurality of policies; and the method further comprises an act of: activating the at least one hardware logic block in response to at least one corresponding policy being loaded.

12. The method of claim 7, wherein: the at least one hardware logic block is configured to construct, based on the metadata input, respective inputs for a plurality of rule queries.
13. The method of claim 12, wherein: the instruction comprises a vector instruction having a plurality of component operations; and the plurality of rule queries correspond, respectively, to the plurality of component operations.
14. The method of claim 12, wherein: the instruction comprises a direct memory access instruction having at least one load component and at least one store component; and the plurality of rule queries comprise at least first and second queries corresponding, respectively, to the at least one load component and the at least one store component.
15. The method of claim 1, wherein: the application data storage location comprises a register or memory location used by the instruction; and the one or more classification bits indicate the metadata input is for determining whether the instruction is allowed to use the register or memory location.
16. The method of claim 15, wherein: processing the metadata input based on the one or more classification bits comprises: using the metadata input to construct an input for a rule query to determine whether the metadata input matches one or more allowed input metadata patterns.
17. The method of claim 1, wherein: the application data storage location comprises a first register or memory location used by the instruction; and the one or more classification bits indicate the metadata input is for performing a metadata update for a second register or memory location used by the instruction.
18. The method of claim 17, wherein: processing the metadata input based on the one or more classification bits comprises: using the metadata input to construct an input for a rule query to determine whether the metadata input matches an entry that maps an input metadata pattern to an output metadata pattern.
19. The method of claim 1, wherein: the one or more classification bits comprises one or more first classification bits indicating that a first portion of the metadata input relates to content stored at a location; and the one or more classification bits further comprises one or more second classification bits indicating that a second portion of the metadata input relates to the location at which the content is stored.
20. The method of claim 1, wherein: processing the metadata input based on the one or more classification bits comprises: applying one or more first update semantics to a first portion of the metadata input; and applying one or more second update semantics to a second portion of the metadata input, wherein the one or more second update semantics are different from the one or more first update semantics.
21. The method of claim 1, wherein: the one or more first bit positions of the metadata input comprise a first portion and a second portion; the one or more classification bits comprise one or more first classification bits associated with the first portion; and processing the metadata input based on the one or more classification bits comprises: modifying the first portion based on the one or more first classification bits.
22. A system comprising circuitry configured to perform a method comprising acts of: receiving one or more metadata inputs associated with an instruction, wherein: the one or more metadata inputs comprises a metadata input from a metadata storage location corresponding to an application data storage location used by the instruction; the metadata input comprises one or more first bit positions and one or more second bit positions; the one or more first bit positions of the metadata input encode at least one metadata symbol; and the one or more second bit positions of the metadata input comprise one or more classification bits indicating one or more attribute values associated with the at least one metadata symbol; and processing the one or more metadata inputs to determine whether to allow the instruction, comprising: identifying the one or more classification bits from the metadata input of the one or more metadata inputs; and processing the metadata input based on the one or more classification bits.

23. At least one non-transitory computer-readable storage medium having stored thereon at least one hardware description which, when synthesized, produces at least one netlist for circuitry configured to perform a method comprising acts of: receiving one or more metadata inputs associated with an instruction, wherein: the one or more metadata inputs comprises a metadata input from a metadata storage location corresponding to an application data storage location used by the instruction; the metadata input comprises one or more first bit positions and one or more second bit positions; the one or more first bit positions of the metadata input encode at least one metadata symbol; and the one or more second bit positions of the metadata input comprise one or more classification bits indicating one or more attribute values associated with the at least one metadata symbol; and processing the one or more metadata inputs to determine whether to allow the instruction, comprising: identifying the one or more classification bits from the metadata input of the one or more metadata inputs; and processing the metadata input based on the one or more classification bits.
